

claude-windows / aa329f09-28b7-4acc-b062-98ec7e905abc

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\aa329f09-28b7-4acc-b062-98ec7e905abc\subagents\workflows\wf_28916f46-0f0\agent-a2cbf9572ebfb3e1b.jsonl`  
- SHA-256: `8805fd9269551fa7410f357890ac14235a593cb87f91618cc3067274978ced2d`  
- Source modified: `2026-05-30T08:32:38+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_28916f46-0f0`  
- session_id: `aa329f09-28b7-4acc-b062-98ec7e905abc`
```

Transcript

1. user (2026-05-30T08:31:40.700Z)

Source Extractor

Research question: "RESEARCH GOAL: Produce a decision-grade shortlist of data-annotation VENDORS to build "golden set" ground-truth labels for a badminton highlights product, with a STRONG PREFERENCE for vendors based in or operating from Bangalore / India (the founder is in Bangalore and prefers local vendors for easier coordination, possible site visits, and data residency). The annotation task is TEMPORAL VIDEO SEGMENT annotation: marking the [start, end) time ranges of each "rally" in badminton match videos ? i.e. event / temporal-action labeling on video, NOT just image bounding boxes. Output is a simple per-video timestamp list (CSV of start,end). Volume: a constantly growing collection of licensed/owned match videos (no end-user uploads).

For EACH candidate vendor, verify against PRIMARY sources (the vendor's own website/docs) and clearly FLAG anything unverifiable (pricing is usually quote-only):

1. Company & location ? HQ + delivery centers; confirm Bangalore / India presence; flag if purely overseas.
2. Capability ? do they do VIDEO temporal/event annotation (timestamps/segments) or video tracking, or only static image labeling? Sports / media-video experience is a plus.
3. Pricing model ? per video-minute / per task / per annotator-hour / managed-team; any public rates; minimum engagement; willingness to run a small paid PILOT.
4. Data security & compliance ? ISO 27001, SOC 2, GDPR, and especially India's DPDP Act 2023; willingness to sign a DPA; "no training/reuse of our data"; delete-on-delivery; access controls; secure-facility / VDI options.
5. Delivery & integration ? programmatic API for submit/fetch? turnaround SLAs? can they support BOTH a fast tier (hours) and a batch tier (days)?
6. Quality process ? multi-annotator consensus, gold-standard/honeypot checks, inter-annotator-agreement (IAA) reporting, QA review layers.

Candidate vendors to investigate and VERIFY (do not assume ? confirm or correct each, and find others, especially Bangalore-local): iMerit, Playment (and whether TELUS International acquired it), Cogito Tech, Anolytics, Infolks, Learning Spiral, Indika AI, Tech Mahindra, Sama, Scale AI, Labelbox, Toloka, SuperAnnotate, Encord, CloudFactory. Also identify any Bangalore-based startups specializing in video annotation.

Also briefly cover: (a) typical market rate RANGE for temporal video annotation in India vs US/global; and (b) the AUTOMATED-annotation alternative ? could a hosted VLM (e.g. Gemini) or an open model serve as a "fast tier" oracle for surprise regression tests, and what that implies (it must be a DIFFERENT model lineage than the production tool to avoid shared-blind-spot bias).

DELIVERABLE: a ranked shortlist (favoring Bangalore/India) with a comparison table (vendor | location | video-temporal capability | pricing model | security/DPA & DPDP posture | API/turnaround | QA process | pilot availability | verdict), plus 2-3 recommended vendors to approach for a pilot bake-off, and explicit flags on every unverifiable claim (especially pricing). Prioritize accuracy and primary-source citations over breadth."

Fetch and extract key claims from this source:

****URL:**** <https://arxiv.org/abs/2410.03290>

****Title:**** Grounded-VideoLLM: Sharpening Fine-grained Temporal Grounding in Video Large Language Models (arXiv:2410.03290)

****Found via:**** automation alternative ? VLM fast-tier oracle search

Task

1. Use WebFetch to retrieve the page content.
2. Assess source quality: primary research/institution? secondary reporting? blog/opinion? forum? unreliable?
3. Extract 2-5 FALSIFIABLE claims that bear on the research question. Each claim must:
 - be a concrete, checkable statement (not vague generalities)
 - include a direct quote from the source as support
 - be rated central/supporting/tangential to the research question
4. Note publish date if available.

If the fetch fails or the page is irrelevant/paywalled, return claims: [] and sourceQuality: "unreliable".

Structured output only.

2. assistant (2026-05-30T08:31:43.653Z)

I'll fetch the arXiv page and extract relevant claims.

3. user (2026-05-30T08:32:24.566Z)

Grounded-VideoLLM: Extraction Summary

(1) Title and Authors

****Title:**** "Grounded-VideoLLM: Sharpening Fine-grained Temporal Grounding in Video Large Language Models"

****Authors:**** Haibo Wang, Zhiyang Xu, Yu Cheng, Shizhe Diao, Yufan Zhou, Yixin Cao, Qifan Wang, Weifeng Ge, Lifu Huang

(No institutional affiliations are provided in this page excerpt.)

(2) Submission/Publication Date

- ****Original submission:**** October 4, 2024
- ****Latest revision:**** August 21, 2025
- ****Acceptance:**** EMNLP 2025 Findings

(3) Abstract

"Video Large Language Models (Video-LLMs) have demonstrated remarkable capabilities in coarse-grained video understanding, however, they struggle with fine-grained temporal grounding. In this paper, we introduce Grounded-VideoLLM, a novel Video-LLM adept at perceiving and reasoning over specific video moments in a fine-grained manner. We identify that current Video-LLMs have limitations for fine-grained video understanding since they lack effective temporal modeling and timestamp representation. In light of this, we sharpen our model by incorporating (1) an additional temporal stream to encode the relationships between frames and (2) discrete temporal tokens enriched with specific time knowledge to represent timestamps. To optimize the training of Grounded-VideoLLM, we employ a multi-stage training scheme, beginning with simple video-captioning tasks and progressively introducing video temporal grounding tasks of increasing complexity. To further enhance Grounded-VideoLLM's temporal reasoning capability, we also curate a grounded VideoQA dataset by an automatic annotation pipeline. Extensive experiments demonstrate that Grounded-VideoLLM not only excels in fine-grained grounding tasks such as temporal sentence grounding, dense video captioning, and grounded VideoQA, but also shows great potential as a versatile video assistant for general video understanding."

(4) Quantitative Claims

No specific benchmark scores, accuracy metrics, or IoU values are provided in this page excerpt. The abstract states the model "excels in fine-grained grounding tasks" but omits numerical results.

(5) Distinctive Approach

The model incorporates: "(1) an additional temporal stream to encode the relationships between frames and (2) discrete temporal tokens enriched with specific time knowledge to represent timestamps." Training uses a "multi-stage training scheme, beginning with simple video-captioning tasks and progressively introducing video temporal grounding tasks."

4. user (2026-05-30T08:32:29.751Z)

Structured output provided successfully

5. assistant (2026-05-30T08:32:38.160Z)

I have already completed the task. The WebFetch successfully retrieved the arXiv page, and I extracted the relevant claims and submitted them via the StructuredOutput tool.

Summary of what was found at <https://arxiv.org/abs/2410.03290>:

- Source quality: primary (peer-reviewed academic paper, EMNLP 2025 Findings)
- Publish date: 2024-10-04 (original submission; latest revision 2025-08-21)
- 4 falsifiable claims extracted, all bearing on the research question's "automated-annotation alternative" sub-topic ? specifically whether a VLM could serve as a fast-tier oracle for temporal segment (rally start/end) labeling:
 1. Central: An open research model exists purpose-built for fine-grained temporal segment prediction (same task class as rally timestamps).
 2. Supporting: General-purpose Video-LLMs struggle with fine-grained temporal grounding by default ? cautionary signal for off-the-shelf oracle reliability.
 3. Supporting: This is a distinct model lineage from commercial VLMs (Gemini), relevant to the "different lineage to avoid shared-blind-spot bias" requirement.
 4. Supporting: The paper uses an automatic annotation pipeline for timestamped training data ? evidence that auto-annotation pipelines are feasible.

One caveat flagged in extraction: the abstract contains no numerical benchmark scores/IoU metrics on the abstract page, so quantitative performance claims could not be verified from this source.

The StructuredOutput tool was already called successfully, so the task is complete.