

claude-windows / aa329f09-28b7-4acc-b062-98ec7e905abc

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\aa329f09-28b7-4acc-b062-98ec7e905abc\subagents\workflows\wf_ddd29371-8c2\agent-a9a2f90233dc4fc84.jsonl`  
- SHA-256: `dc7b32484c660347db3f247d89804a2a3099cc438d645ad3cf0c3fc9f6ae6c3e`  
- Source modified: `2026-05-30T03:12:45+00:00`  
- Imported at: `2026-07-05T16:48:25+00:00`  
- project: `wf_ddd29371-8c2`  
- session_id: `aa329f09-28b7-4acc-b062-98ec7e905abc`
```

Transcript

1. user (2026-05-30T03:11:57.138Z)

Adversarial Claim Verifier (voter 1/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

RESEARCH GOAL: Identify and eliminate every dependency that could legally or economically hinder monetizing the "badminton-highlight-indexer" project as a HOSTED API SaaS. Business model: the project runs server-side on a cloud GPU VM (NVIDIA); users upload their own badminton videos and receive auto-generated highlight reels. CRITICAL CONSTRAINT: we do NOT distribute the code or any binary to users ? everything stays server-side ? so software that is merely *hosted* on the server (invoked as a separate process, e.g. ffmpeg) is acceptable where distribution-triggered copyleft (GPL/LGPL) would NOT be triggered, BUT network/SaaS copyleft (AGPL §13) and external metered-API terms STILL apply. Preference order: permissively-licensed (Apache-2.0/MIT/BSD) and royalty-free, self-hostable components. Paid commercial licenses (e.g. Ultralytics Enterprise, codec patent licenses) are acceptable ONLY as clearly-flagged fallbacks. Output must be decision-grade with licenses verified against PRIMARY sources (actual GitHub LICENSE files, official pricing/terms pages), because this drives a real business decision ? flag any claim you could not verify.

CURRENT STACK TO AUDIT (verify each license + its SaaS implication):

- ultralytics (Python lib) + the bundled YOLOv8n.pt model weights ? believed AGPL-3.0. Confirm the license of BOTH the code and the pretrained weights, whether the AGPL network clause is triggered by a closed-source hosted service that uses ultralytics server-side, and the existence/cost/terms of the Ultralytics Enterprise license.
- google-genai SDK + the Google Gemini API (model gemini-2.5-flash). Distinguish the SDK license (Apache?) from the SERVICE terms. Confirm: (a) is commercial use of outputs allowed in a paid product; (b) does Google train on / human-review prompt+video data on the PAID/billed tier vs the free tier; (c) abuse-monitoring / data-retention terms; (d) current per-token / per-video pricing for gemini-2.5-flash (input video tokens-per-second, output) so we can sketch per-video cost.
- ffmpeg / ffprobe (system binaries, invoked as separate processes). Clarify GPL vs LGPL builds and whether pure server-side SaaS use (no distribution to users) triggers any source-disclosure obligation. SEPARATELY and importantly: video CODEC PATENT royalties for a commercial service that DECODES user-uploaded H.264/AVC and H.265/HEVC (e.g. GoPro footage) and ENCODES output ? covering MPEG-LA / Via-LA (AVC), Access Advance + Via-LA + others (HEVC) pools, AAC audio (Via-LA), what activities trigger royalties, any free thresholds (e.g. the AVC sub-100k or free-internet-broadcast provisions), and whether a small SaaS realistically incurs these. Then royalty-FREE codec alternatives for the OUTPUT we generate: AV1, VP9, Opus ? and their encoder maturity/speed on a GPU VM.
- torch, torchvision ? confirm BSD; flag that some torchvision *pretrained weights* may carry separate licenses.
- opencv-python ? confirm Apache-2.0 and that the PyPI wheel's bundled components are clean for commercial server use.

- lapx (a fork of "lap" for ByteTrack association) ? confirm license.
- Quick permissive confirmation for: fastapi, uvicorn, pydantic, jinja2, python-dotenv, numpy.

PERMISSIVE ALTERNATIVES TO RESEARCH (all must be GPU-VM-friendly and usable in a closed-source commercial SaaS):

1. Object detection to replace ultralytics/YOLOv8 for player/person detection (and possibly shuttle): compare YOLOX (Apache-2.0), RT-DETR / RT-DETRv2 (the Baidu/Apache implementations ? NOT the Ultralytics port), PP-YOLOE, MMDetection (Apache-2.0), Detectron2 (Apache-2.0), torchvision built-in detectors (Faster R-CNN/RetinaNet/FCOS/SSD, BSD), RF-DETR (Roboflow), and D-FINE/DEIM. For EACH, verify BOTH the code license AND the pretrained-weights license (some "open" detectors like YOLO-NAS have restrictive weight licenses ? confirm). Note accuracy/speed trade-offs for person detection on sports video.
2. Shuttle/ball trajectory tracking ? VERIFY the actual repo licenses (commercial-use viability) of: WASB / WASB-SBDT ("Widely Applicable Strong Baseline" for sports ball detection), TrackNetV2 and TrackNetV3, MonoTrack, and any permissively-licensed alternative for tiny-fast-object tracking. This is the planned shuttle-tracking substrate, so its commercial-use license is high-stakes.
3. LLM / semantic rally-boundary step ? research ALL THREE strategies (the user wants all three compared):
 - (a) ELIMINATE the hosted LLM entirely: permissively-licensed temporal action segmentation / video models suitable for learning rally boundaries offline (e.g. MS-TCN/MS-TCN++, ASFormer, ActionFormer, and similar) ? verify licenses and whether they fit a from-annotations learned segmenter.
 - (b) SELF-HOST an open-weights vision-language model that can ingest video frames for semantic boundaries: compare Qwen2.5-VL, InternVL2/2.5, LLaVA-OneVision/LLaVA-Video, MiniCPM-V, VideoLLaMA ? for EACH verify the WEIGHTS license for commercial use (Apache-2.0 vs custom "research-only"/"non-commercial"/community licenses with MAU caps) and the GPU VRAM needed per size tier.
 - (c) KEEP Gemini but de-risk: summarize commercial terms, data-usage on paid tier, rate limits, and per-video cost at scale, plus other hosted-API alternatives (OpenAI, Anthropic) and how their commercial/data terms compare.

DELIVERABLE: A cited report containing (1) a per-dependency RISK TABLE ? columns: dependency, license, does-SaaS-hosting-trigger-obligations?, monetization verdict (CLEAR / WATCH / BLOCKER), recommended action; (2) a recommended FULLY-PERMISSIVE target stack that eliminates all BLOCKERS; (3) paid-license fallbacks where relevant; (4) a rough per-video processing-cost sketch on a GPU VM for each of the three LLM strategies; (5) an explicit list of the codec/patent considerations and the royalty-free output-codec recommendation. Prioritize accuracy and primary-source citations over breadth.

Claim under review

"On Paid Services (including the paid quota of the Gemini API), Google does not use prompts (including system instructions, cached content, and files such as images, videos, or documents) or responses to improve its products. This is the key data-usage assurance for a commercial SaaS billed via the paid Gemini API."

****Source:**** <https://ai.google.dev/gemini-api/docs/data-usage> (primary)

****Supporting quote:**** "When you use Paid Services, including, for example, the paid quota of the Gemini API, Google doesn't use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve our products"

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-05-30T03:11:59.764Z)

I'll verify this claim about Google's data usage policy for Paid Services on the Gemini API. Let me start by loading the web tools and fetching the primary source.

3. user (2026-05-30T03:12:18.181Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

4. user (2026-05-30T03:12:21.346Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

5. user (2026-05-30T03:12:23.600Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

6. user (2026-05-30T03:12:26.788Z)

Web search results for query: "Gemini API paid tier Google does not use prompts to improve products data usage exceptions abuse monitoring"

Links: [{"title":"Billing | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/billing"}, {"title":"Gemini Developer API pricing | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/pricing"}, {"title":"Is Gemini API Free? (2025) - Spur","url":"https://www.spurnow.com/en/blogs/is-gemini-api-free"}, {"title":"Gemini API Pricing Calculator & Cost Guide (May 2026)","url":"https://costgoat.com/pricing/gemini-api"}, {"title":"Gemini API Pricing 2026: Complete Cost Guide for All Models | metactoo","url":"https://www.metactoo.com/blogs/the-true-cost-of-google-gemini-a-guide-to-api-pricing-and-integration"}, {"title":"Gemini API Additional Terms of Service | Google AI for Developers","url":"https://ai.google.dev/gemini-api/terms"}, {"title":"Gemini Pricing in 2026 for Individuals, Orgs & Developers","url":"https://www.finout.io/blog/gemini-pricing-in-2026"}, {"title":"New Google AI Studio Terms Ban Consumer Use: The Developer Pivot","url":"https://www.remio.ai/post/new-google-ai-studio-terms-ban-consumer-use-the-developer-pivot"}, {"title":"Gemini API Key Free in 2026: What Is Free, Where Limits Live, and Why More Keys Do Not Add Quota | LaoZhang AI Blog","url":"https://blog.laozhang.ai/en/posts/gemini-api-free-tier"}, {"title":"How Gemini for Google Cloud uses your data | Google Cloud Documentation","url":"https://docs.cloud.google.com/gemini/docs/discover/data-governance"}]

Based on the web search results, here's what I found regarding Gemini API paid tier and data usage:

Data Usage on Paid Tiers

When you link a billing account and move to the Paid Tiers, Google ensures your prompts and responses are not used to improve Google products. On Paid Services, Google does not use your prompts or responses to improve their models, and they're processed under a data processing addendum with logging mainly for abuse detection and compliance.

Free vs. Paid Tier Differences

The key distinction between tiers is clear: On the free tier, prompts and responses can be used (with human review) to improve Google's products. When you use Unpaid Services, Google uses the content you submit and any generated responses to provide, improve, and develop Google products and services. To help with quality and improve products, human reviewers may read, annotate, and process your API input and output, though Google takes steps to protect your privacy as part of this process.

Data Processing Addendum

The Terms of service details how data is handled when Cloud billing is enabled, and your Google AI Studio prompts are treated under the same "Paid Services" terms so long as at least 1 API project has billing enabled.

In summary, the main exceptions to Google using your data for product improvement are when you have an active paid tier billing account, which then subjects your data to a data processing addendum with logging limited to abuse detection and compliance purposes.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

7. user (2026-05-30T03:12:30.262Z)

Data Usage Comparison: Paid vs Unpaid Services

Unpaid Services

Google **does** use prompts and responses to improve products:

> "Google uses the content you submit to the Services and any generated responses to provide, improve, and develop Google products and services"

Additionally: "Human reviewers may read, annotate, and process your API input and output."

Paid Services

Google **does not** use prompts/responses for product improvement:

> "Google doesn't use your prompts...or responses to improve our products"

Exceptions for Paid Services (data retention/monitoring):

1. **Abuse Prevention**: Data is "logged...solely for detecting and preventing violations of the Prohibited Use Policy"
2. **Security**: Retention occurs "to maintain the safety and security of the Services"
3. **Legal Compliance**: Data may be retained for "any required legal or regulatory disclosures"
4. **Storage Duration**: This data "may be stored transiently or cached in any country in which Google or its agents maintain facilities"

Additional Data Categories

Other operational information?including account details, billing history, usage metrics, and technical performance data?remains governed by Google's standard privacy policies regardless of service tier.

8. user (2026-05-30T03:12:29.542Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

9. user (2026-05-30T03:12:31.432Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

10. user (2026-05-30T03:12:35.737Z)

Web search results for query: "Gemini API paid services data usage Google does not use prompts improve products"

Links: [{"title": "Gemini API Additional Terms of Service | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/terms"}, {"title": "Billing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/billing"}, {"title": "How Gemini for Google Cloud uses your data | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini/docs/discover/data-governance"}, {"title": "Zero data retention in the Gemini Developer API | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/zdr"}, {"title": "Clarify use of user's code for training within the free quota on paid service · Issue #1472 · google-gemini/gemini-cli", "url": "https://github.com/google-gemini/gemini-cli/issues/1472"}, {"title": "Data Logging and Sharing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/logs-policy"}, {"title": "Is Gemini API Free? (2025) - Spur", "url": "https://www.spurnow.com/en/blogs/is-gemini-api-free"}, {"title": "Generative AI APIs Additional Terms of Service | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/terms-archive/terms_12_20_24"}]

Based on the search results, here's what I found regarding Gemini API paid services and data usage:

Data Usage Policy for Paid Services

When you use Paid Services with Gemini API, Google doesn't use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve our products.

Key Distinctions Between Free and Paid Tiers

To access higher rate limits, use advanced models, and ensure your prompts and responses are not used to improve Google products, you can link a billing account and prepay to move to the Paid Tiers.

In contrast, when you use Unpaid Services, including Google AI Studio and the unpaid quota on Gemini API, Google uses the content you submit to the Services and any generated responses to provide, improve, and develop Google products and services.

Limited Data Retention

For Paid Services, Google logs prompts and responses for a limited period of time, solely for the purpose of detecting violations of the Prohibited Use Policy and any required legal or regulatory disclosures.

The fundamental distinction is clear: with paid Gemini API services, your data is protected from being used for Google product improvement, while free tier usage allows Google to use your submissions for product development and improvement.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

11. user (2026-05-30T03:12:36.857Z)

Web search results for query: "Gemini API paid services Google does not use prompts to improve

products data usage abuse monitoring human review"

Links: [{"title": "Abuse monitoring | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/usage-policies"}, {"title": "Gemini Apps Privacy Hub - Gemini Apps Help", "url": "https://support.google.com/gemini/answer/13594961?hl=en"}, {"title": "Gemini API Additional Terms of Service | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/terms_preview"}, {"title": "Is Gemini API Free? (2025) - Spurn", "url": "https://www.spurnow.com/en/blogs/is-gemini-api-free"}, {"title": "Billing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/billing"}, {"title": "Is Gemini Safe for Legal Work? Protecting Confidential Information - Spellbook", "url": "https://spellbook.com/learn/is-gemini-private"}, {"title": "Google Gemini Privacy Policy Explained | Cape - Cape", "url": "https://www.cape.co/blog/google-gemini-privacy-policy"}, {"title": "How Gemini for Google Cloud uses your data | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini/docs/discover/data-governance"}, {"title": "Monitor Gemini Code Assist usage | Gemini for Google Cloud | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini/docs/codeassist/monitor-gemini-code-assist"}, {"title": "Review Gemini usage in your organization | Google Workspace with Gemini | Google Workspace Help", "url": "https://knowledge.workspace.google.com/admin/gemini/review-gemini-usage-in-your-organization"}]

Based on the search results, here's what I found regarding Gemini API paid services and data usage:

Gemini API Paid Services - Data Usage

When you use Paid Services (such as the paid quota of the Gemini API), Google doesn't use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve our products.

Logging and Abuse Monitoring

For Paid Services, Google logs prompts and responses for a limited period of time, solely for the purpose of detecting violations of the Prohibited Use Policy and any required legal or regulatory disclosures.

Data can be accessed for human review only by authorized Google employees via an internal governance assessment and review management platform.

Key Differences: Paid vs. Unpaid Services

For comparison, when you use Unpaid Services (including Google AI Studio and the unpaid quota on Gemini API), Google uses the content you submit to the Services and any generated responses to provide, improve, and develop Google products and services and machine learning technologies, and human reviewers may read, annotate, and process your API input and output.

The primary distinction is that paid Gemini API services provide stronger privacy protections by preventing Google from using your data for product improvement, while unpaid services allow such usage with human review involved.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

12. user (2026-05-30T03:12:36.479Z)

Web search results for query: ""Gemini API" paid tier data usage "doesn't use your prompts" abuse monitoring human reviewers retention"

Links: [{"title": "Abuse monitoring | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/usage-policies"}, {"title": "Gemini Apps Privacy Hub - Gemini Apps Help", "url": "https://support.google.com/gemini/answer/13594961?hl=en"}, {"title": "Gemini API and"}]

Your Data Privacy: A 2025 Guide for Privacy-Conscious Users", "url": "https://redact.dev/blog/gemini-api-terms-2025"}, {"title": "Everything You Should Know About Google Gemini Data Retention Policy", "url": "https://anarlog.so/blog/google-gemini-data-retention-policy/"}, {"title": "Safety and factuality guidance | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/safety-guidance"}, {"title": "How Gemini for Google Cloud uses your data | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini/docs/discover/data-governance"}, {"title": "What Really Happens to Your Gen AI Data - Cerium Networks", "url": "https://ceriumnetworks.com/what-really-happens-to-your-gen-ai-data/"}, {"title": "Gemini API Additional Terms of Service | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/terms"}, {"title": "Generative AI in Google Workspace Privacy Hub | Google Workspace with Gemini | Google Workspace Help", "url": "https://knowledge.workspace.google.com/admin/gemini/generative-ai-in-google-workspace-privacy-hub"}, {"title": "Is Gemini API Free? (2025) - Spur", "url": "https://www.spurnow.com/en/blogs/is-gemini-api-free"}]

Based on the search results, here's what I found about the Gemini API paid tier data usage, the "doesn't use your prompts" policy, abuse monitoring, and human reviewers:

Data Usage on Paid Tiers

On paid tiers, prompts and responses are not used to improve Google's models and are processed under Google's data processing addendum. More specifically, Gemini doesn't use your prompts or its responses as data to train its models.

Data Retention for Abuse Monitoring

Google retains API inputs and outputs for up to 55 days for abuse monitoring. This data is not used for model training.

Human Review and Monitoring

The human review process for paid tiers differs from free tiers. Google retains prompts, contextual information, and outputs for 55 days to detect and prevent violations of the Prohibited Use Policy. When prompts or model outputs are flagged by safety filters and abuse detection systems, authorized Google employees may assess the flagged content and either confirm or correct the classification based on predefined guidelines and policies.

Data can be accessed for human review only by authorized Google employees via an internal governance assessment and review management platform.

For enterprise/paid services specifically, Google does not use your prompts or outputs to train models or improve services without consent.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

13. user (2026-05-30T03:12:37.554Z)

Web search results for query: "Gemini API paid tier Google data usage policy prompts responses improve products"

Links: [{"title": "Gemini API Additional Terms of Service | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/terms"}, {"title": "Billing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/billing"}, {"title": "How Gemini for Google Cloud uses your data | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini/docs/discover/data-governance"}, {"title": "Is Gemini API Free? (2025) - Spur", "url": "https://www.spurnow.com/en/blogs/is-gemini-api-free"}, {"title": "Gemini API Pricing Calculator & Cost Guide (May 2026)", "url": "https://costgoat.com/pricing/gemini-api"}, {"title": "Gemini Developer API pricing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/pricing"}, {"title": "Clarify use of"}]

```
user's code for training within the free quota on paid service · Issue #1472 · google-gemini/gemini-cli", "url": "https://github.com/google-gemini/gemini-cli/issues/1472"}, {"title": "New Google AI Studio Terms Ban Consumer Use: The Developer Pivot", "url": "https://www.remio.ai/post/new-google-ai-studio-terms-ban-consumer-use-the-developer-pivot"}, {"title": "Gemini API Free Tier: Complete Guide to Rate Limits, Models & Getting Started (2026) | YingTu", "url": "https://yingtu.ai/en/blog/gemini-api-free-tier"}, {"title": "Gemini API Free Tier 2026: 1,500 Req/Day, 1M TPM ? No Card - TokenMix Blog", "url": "https://tokenmix.ai/blog/gemini-api-free-tier-limits"}]]
```

I found comprehensive information about Google's data usage policies for the Gemini API paid tier. Here are the key findings:

Data Usage Policy for Paid Tier

When you use Paid Services on Gemini API, Google doesn't use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve products.

Contrast with Unpaid Services

When you use Unpaid Services, Google uses the content you submit and any generated responses to provide, improve, and develop Google products and services and machine learning technologies. To help with quality and improve products, human reviewers may read, annotate, and process your API input and output.

Data Handling for Paid Services

For Paid Services, Google logs prompts and responses for a limited period of time, solely for detecting and preventing violations of the Prohibited Use Policy to maintain the safety and security of the Services, and any required legal or regulatory disclosures.

Key Benefit

To access higher rate limits, use advanced models, and ensure your prompts and responses are not used to improve Google products, you can link a billing account and prepay to move to the Paid Tiers.

The distinction is clear: the paid tier provides a privacy guarantee that your data won't be used for product improvement, whereas the free tier explicitly allows Google to use your data in this way.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

14. user (2026-05-30T03:12:39.730Z)

Tool permission request failed: Error: Tool permission stream closed before response received

15. user (2026-05-30T03:12:39.730Z)

```
<tool_use_error>Cancelled: parallel tool call Bash(curl -sL "https://ai.google.dev/gemini-a?)"
errored</tool_use_error>
```

16. user (2026-05-30T03:12:39.730Z)

```
<tool_use_error>Cancelled: parallel tool call Bash(curl -sL "https://ai.google.dev/gemini-a?)"
errored</tool_use_error>
```

17. user (2026-05-30T03:12:40.042Z)

```
<tool_use_error>Cancelled: parallel tool call Bash(curl -sL "https://ai.google.dev/gemini-a?)"
```


errored</tool_use_error>

18. user (2026-05-30T03:12:42.296Z)

<tool_use_error>Cancelled: parallel tool call Bash(curl -sL "https://ai.google.dev/gemini-a?")
errored</tool_use_error>

19. user (2026-05-30T03:12:44.236Z)

<tool_use_error>Cancelled: parallel tool call Bash(curl -sL "https://ai.google.dev/gemini-a?")
errored</tool_use_error>

20. user (2026-05-30T03:12:45.241Z)

<tool_use_error>Cancelled: parallel tool call Bash(curl -sL "https://ai.google.dev/gemini-a?")
errored</tool_use_error>