

claude-windows / aa329f09-28b7-4acc-b062-98ec7e905abc

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\aa329f09-28b7-4acc-b062-98ec7e905abc\subagents\workflows\wf_28916f46-0f0\agent-a83e4a4128377c498.jsonl`
- SHA-256: `a33492250af2b4f81d8573c01f3c8bc2dd4449176440f637b6184e4e50ae600b`
- Source modified: `2026-05-30T09:41:07+00:00`
- Imported at: `2026-07-05T16:48:24+00:00`
- project: `wf_28916f46-0f0`
- session_id: `aa329f09-28b7-4acc-b062-98ec7e905abc`
```

Transcript

1. user (2026-05-30T09:41:06.015Z)

Synthesis: research report

Question: RESEARCH GOAL: Produce a decision-grade shortlist of data-annotation VENDORS to build "golden set" ground-truth labels for a badminton highlights product, with a STRONG PREFERENCE for vendors based in or operating from Bangalore / India (the founder is in Bangalore and prefers local vendors for easier coordination, possible site visits, and data residency). The annotation task is TEMPORAL VIDEO SEGMENT annotation: marking the [start, end) time ranges of each "rally" in badminton match videos ? i.e. event / temporal-action labeling on video, NOT just image bounding boxes. Output is a simple per-video timestamp list (CSV of start,end). Volume: a constantly growing collection of licensed/owned match videos (no end-user uploads).

For EACH candidate vendor, verify against PRIMARY sources (the vendor's own website/docs) and clearly FLAG anything unverifiable (pricing is usually quote-only):

1. Company & location ? HQ + delivery centers; confirm Bangalore / India presence; flag if purely overseas.
2. Capability ? do they do VIDEO temporal/event annotation (timestamps/segments) or video tracking, or only static image labeling? Sports / media-video experience is a plus.
3. Pricing model ? per video-minute / per task / per annotator-hour / managed-team; any public rates; minimum engagement; willingness to run a small paid PILOT.
4. Data security & compliance ? ISO 27001, SOC 2, GDPR, and especially India's DPDP Act 2023; willingness to sign a DPA; "no training/reuse of our data"; delete-on-delivery; access controls; secure-facility / VDI options.
5. Delivery & integration ? programmatic API for submit/fetch? turnaround SLAs? can they support BOTH a fast tier (hours) and a batch tier (days)?
6. Quality process ? multi-annotator consensus, gold-standard/honeypot checks, inter-annotator-agreement (IAA) reporting, QA review layers.

Candidate vendors to investigate and VERIFY (do not assume ? confirm or correct each, and find others, especially Bangalore-local): iMerit, Playment (and whether TELUS International acquired it), Cogito Tech, Anolytics, Infolks, Learning Spiral, Indika AI, Tech Mahindra, Sama, Scale AI, Labelbox, Toloka, SuperAnnotate, Encord, CloudFactory. Also identify any Bangalore-based startups specializing in video annotation.

Also briefly cover: (a) typical market rate RANGE for temporal video annotation in India vs US/global; and (b) the AUTOMATED-annotation alternative ? could a hosted VLM (e.g. Gemini) or an open model serve as a "fast tier" oracle for surprise regression tests, and what that implies (it must be a DIFFERENT model lineage than the production tool to avoid shared-blind-spot bias).

DELIVERABLE: a ranked shortlist (favoring Bangalore/India) with a comparison table (vendor | location | video-temporal capability | pricing model | security/DPA & DPDP posture | API/turnaround | QA process | pilot availability | verdict), plus 2-3 recommended vendors to approach for a pilot bake-off, and explicit flags on every unverifiable claim (especially pricing). Prioritize accuracy and primary-source citations over breadth.

15 claims survived 3-vote adversarial verification. Merge semantic duplicates and synthesize.

Confirmed claims

[0] iMerit's dedicated video-annotation page frames its capability around per-frame object-tracking techniques (bounding boxes, 3D c

Vote: 3-0 · Source: <https://imerit.net/computer-vision/video-annotation/> (primary)

Quote: "iMerit delivers video annotation solutions with advanced labeling techniques, object tracking, and quality control to improve model accuracy and performance."

Verifier evidence (high): VERIFIED against iMerit's own live content via a clean WebFetch of the authoritative page (the cited <https://imerit.net/computer-vision/video-annotation/> returns 301 -> imerit.ai; the equivalent live content page <https://imerit.ai/solutions/computer-vision/video-annotation-services/> was fetched cleanly and is the canonical version of this material). Every component of the claim checks out:

- 1) Technique enumeration matches verbatim. The page names exactly: Bounding boxes, Polygon, Semantic segmentation, Keypoint, Landmark, 3D cuboid, Polyline ? plus "video interpolation" ("iMerit's video annotation platform utilizes video interpolation to rapidly annotate suitable video footage"). This is precisely the object-on-frame tracking toolset the claim describes.
- 2) The 'frame-by-frame' detail is correct: WebFetch confirms "frame-by-frame" is PRESENT on the page.
- 3) The 'temporal' and 'action recognition' absence sub-claims are correct: WebFetch term-checks return "Temporal: Absent" and "Action recognition: Absent." Also Absent: event detection, rally, timestamp, start time, end time, highlight, clip. (This corrects my earlier session error where a Wayback snapshot turned out to be a 404 page and gave invalid counts ? those were discarded; the clean live fetch is authoritative.)
- 4) The load-bearing conclusion is fully supported: the page frames capability around per-frame object tracking and interpolation, NOT explicit free-form [start,end) temporal event/rally segmentation. The intro headline ("iMerit delivers stellar video annotation services that power AI, machine learning, and data operation strategies") and the technique list contain nothing about labeling event/segment time-boundaries. Therefore the exact rally-segment-timestamp task fit is genuinely UNVERIFIED from the page and should be scoped in a pilot, exactly as the claim states. The claim is also appropriately hedged (it says fit is unverified, not that iMerit cannot do it), so it is not an overreach.

The supporting quote provided with the claim is accurate iMerit copy and consistent with the page's object-tracking framing.

ONE minor accuracy note for the synthesizer (does NOT affect the verdict): the cited URL path is stale ? <https://imerit.net/computer-vision/video-annotation/> now 301-redirects to the imerit.ai domain (iMerit rebranded), and the live content lives at imerit.ai/solutions/computer-vision/video-annotation-services/. The citation should be updated to the resolving imerit.ai URL, but the substance is unchanged. Separately, note iMerit DOES use "temporal" in unrelated blog content (e.g., a post on temporal drift in generated video), but not on this services page and only in the sense of frame-to-frame consistency ? which would not change the rally-task conclusion regardless.

[1] Encord supports labeling a frame-range time segment by explicitly setting a start frame and end frame for a classification instance

Vote: 3-0 · Source: <https://encord.com/video-annotation-tool/> (primary)

Quote: "Set the desired start and end frame by clicking the **Set start to current** and **Set end to current** buttons."

Verifier evidence (high): VERIFIED, not refuted. The substantive claim ? Encord lets you set an explicit start-frame and end-frame [start,end) temporal span on a video classification instance ? is directly confirmed by Encord's own primary docs at the exact canonical URL the claimant cites (docs.encord.com/platform-documentation/Annotate/annotate-label-editor/annotate-videos). Multiple INDEPENDENT WebSearch passes returned the verbatim feature text attributed to that URL: "you can set the desired start and end frame by clicking the 'Set start to current' and 'Set end to current' buttons or using the set hotkeys c and v, respectively," plus "'Set start to current' (hotkey C) defines the current frame as the start of the range to which the classification is applied" and "'Set end to current' (hotkey V) defines the current frame as the end of the range." The docs further describe an "Add range" (hotkey N) action and a "This frame" (hotkey X) button, and explain this is for classifications that "appear in multiple, non-contiguous sections of a video (such as when a car enters and leaves the frame)." This exactly

matches the supporting quote ("Set the desired start and end frame by clicking the Set start to current and Set end to current buttons") and confirms the [start,end] primitive needed to mark rally time ranges. The claim is NOT an overreach: it accurately describes a real, documented per-instance frame-range mechanism for classifications. Source quality (vendor's own primary product documentation) matches the claim's strength, and this is a routine annotation feature, not an extraordinary or marketing claim. Currency: the canonical docs page is current/live and is returned as the top result for the feature query. The parenthetical FLAG is also accurate: the original task source URL <https://encord.com/video-annotation-tool/> does not resolve to the cited content (a dead/redirected marketing path), so the claimant correctly substituted the live canonical docs URL (docs.encord.com/.../annotate-videos), which is consistently surfaced by search as the authoritative source for this exact feature. CAVEAT on the dead-link sub-claim: my single direct programmatic fetch of encord.com/video-annotation-tool/ did not return a clean confirmation of a literal 'HTTP 404 / Page not found' body in this session, so the precise status code label is the one element I could not independently re-verify byte-for-byte; however, this does not affect the core, well-supported claim about the start/end-frame classification capability, which is the substance under review.

[2] Encord is a video-native annotation platform: annotators label full-length videos directly without breaking them into individual frames

Vote: 2-1 · Source: <https://encord.com/video-annotation-tool/> (primary)

Quote: "Encord allows annotators to directly annotate on full-length videos without breaking them into frames, which saves time and mitigates the risk of missed frames and bounding box drift."

Verifier evidence (high): The supporting quote is verbatim from the primary source (encord.com/video-annotation-tool/), confirmed via search snippet: "Encord allows annotators to directly annotate on full-length videos without breaking them into frames, which saves time and mitigates the risk of missed frames and bounding box drift caused by mismatched frame rates." A second Encord page (encord.com/automate/video-annotation/) repeats: "eliminating the need to fragment videos into individual frames or chunks," and the tool page states videos "up to 2 hours in length" can be annotated "natively." The claim's "preserving temporal context" is corroborated by Encord's own blog ("preserving temporal context across the sequence"). Crucially, the capability is independently corroborated by NEUTRAL third parties, not just vendor copy: V7 Labs (a competitor) describes the same native-video pattern as a category feature; Reddit r/computervision users state "Encord and SuperAnnotate handle long videos natively without frame extraction. CVAT requires you to split videos into frames first"; G2 2025 reviews praise its "native video annotation capabilities." No source contradicts the capability. The only qualifier found is a performance nit (G2: "occasional performance lag on very large video files") ? not a refutation. Minor caveat that does NOT refute: "without breaking into frames" describes the annotator workflow/UX; under the hood video is still frames with keyframe interpolation (Encord docs) ? but the claim is about workflow and temporal-context preservation, which is accurate. The claim is current (2024-2025 sources), the source authority (vendor docs) is appropriate for a product-capability claim, and it is corroborated beyond the vendor. The product-fit editorializing ("substrate suited to long badminton videos") follows logically from verified 2-hour native support. Note: I could not load the page via WebFetch (site fetch errored twice), so verification rests on direct search snippets from encord.com plus independent corroboration rather than a full-page fetch.

[3] Encord supports temporal/event annotation primitives: frame-level classifications (treating a whole frame as a unit) plus dynamic attributes

Vote: 3-0 · Source: <https://encord.com/video-annotation-tool/> (primary)

Quote: "Keyframe-based attribute propagation allows you to mark frames where dynamic attributes change value as keyframes."

Verifier evidence (high): CONFIRMED by primary docs (docs.encord.com), both halves of the compound claim. Frame-level classifications (whole frame as unit), verbatim: "Classifications are used to annotate the whole frame or image, rather than a specific object within the frame. For example, you could use a classification to indicate whether an image is taken during the day or night." Dynamic attributes changing value over a video, verbatim: "A static attribute remains the same throughout the duration of an object, while a dynamic attribute can change over the course of a video. For example, a car might have a static attribute of color (red)... but a dynamic attribute of state (parked, moving) that changes throughout the video." An earlier docs fetch corroborated the keyframe mechanism (setting a value at a frame creates a keyframe that persists until the next change), consistent with the claim's marketing quote on "keyframe-based attribute propagation." The traffic-light/parked-vs-moving examples ARE "classifying states over time," and frame-level classification over frame ranges is directly applicable to per-frame in-rally vs not-in-rally labeling. Checklist: (1) not an overreach ? the single supporting quote covers only the keyframe/dynamic-attribute half, but the frame-classification half is independently confirmed by docs, so the full claim stands; (2) no contradicting evidence found;

(3) source quality = primary vendor docs, sufficient for this mundane claim; (4) not outdated ? classifications + static/dynamic attributes + keyframes are core longstanding Encord features; (5) not fluff ? corroborated by concrete technical docs. NOT extraordinary. Caveats (non-refuting): could NOT load encord.com/video-annotation-tool/ to re-confirm the exact marketing quote verbatim (HTTP 403 + JS-render proxy failure); and the WebSearch tool was non-functional this session (returned model fallback text, not real results), so no independent adversarial web search was possible ? but primary docs affirmatively support the claim.

[4] Encord Annotate supports temporal video annotation: instances created in the video editor are tracked across a range of frames in

Vote: 3-0 · Source: <https://docs.encord.com/platform-documentation/annotate> (primary)

Quote: "Instantiating objects or classifications in the video editor generates a UUID that uniquely identifies that instance across a range of frames (i.e., in the temporal dimension)."

Verifier evidence (high): Claim is well-supported by primary Encord sources, with no contradicting evidence. The claim is a faithful (near-verbatim) paraphrase of the supporting quote ? not an overreach. Multiple independent primary sources corroborate the UUID-across-frames temporal mechanism: (1) SDK ObjectInstance reference (docs.encord.com/sdk-documentation/sdk-references/ObjectInstance): object_hash = "A unique identifier for the object instance"; set_for_frames() "Places the object instance on the specified frame(s)"; get_annotations() "Returns all annotations for the object instance on all frames it has been placed on." (2) Label-editor docs (annotate-images-videos-dicom): "When you create an object in a video, it is automatically assigned a unique identifier. This identifier remains consistent across all frames in which the object appears" + interpolation between keyframes. (3) Video-annotation-tool page: "Annotate video natively without breaking videos into individual image frames," "Temporal classification: Label actions and events as they occur over time across video frames," "Track objects across frames." Source quality matches claim strength (vendor's own docs + technical SDK reference). Minor terminology note only: SDK names the identifier object_hash whereas the doc page calls it a UUID ? same concept (unique per-instance ID persisting across a frame range), not a substantive discrepancy. NOTE for the research goal: this confirms Encord does VIDEO temporal/object annotation generally; it does NOT by itself prove the specific "rally [start,end) timestamp CSV" workflow, though temporal/sequence classification ("label actions and events as they occur over time") is directly applicable.

[5] Encord supports range-level (time-range) classification on video ? annotators set a start and end frame for a classification, which

Vote: 3-0 · Source: <https://docs.encord.com/platform-documentation/annotate> (primary)

Quote: "set the desired start and end frame by clicking the Set start to current and Set end to current buttons or using the set hotkeys c and v, respectively"

Verifier evidence (high): The claim accurately restates Encord's documented capability. Multiple independent WebSearch results consistently quote Encord's own docs (docs.encord.com/platform-documentation/.../annotate-videos and .../gettingstarted-labeling) describing frame-level CLASSIFICATION with start/end frame ranges: "Creating frame classification(s) requires an Ontology with a classification annotation type... set the desired start and end frame by clicking the Set start to current and Set end to current buttons or using the set hotkeys c and v, respectively." This resolves the only material doubt about the claim ? the start/end-frame mechanism is explicitly attached to CLASSIFICATION instances (range-level event labeling), not merely object-track keyframes. Corroborating quotes: "instantiating classifications in the video editor generates a UUID that uniquely identifies that instance across a range of frames (in the temporal dimension)" and "Encord has temporal labeling built into its UI, allowing users to annotate events, transitions or behaviors." Source is Encord's primary documentation, appropriate for a product-capability claim; the platform is current (2026 docs), so no staleness concern; and no contradicting evidence was found across 10+ searches. NOTE: The exact URL cited in the claim (.../platform-documentation/annotate) returns HTTP 404 ? the live path is .../platform-documentation/Annotate/annotate-label-editor/annotate-videos ? but the substance is fully confirmed on the correct page, so this is a citation-URL nit, not a refutation. Direct WebFetch/curl to docs.encord.com returned 404/empty in this environment (likely JS-rendered or path-changed), but WebSearch snippets quote the primary docs verbatim and agree unanimously. The supporting quote genuinely supports the claim, source quality matches claim strength, and the claim is current and accurate.

[6] Sama DOES perform temporal video annotation suitable for marking time-based events: its site lists 'Object Tracking and Temporal

Vote: 3-0 · Source: <https://www.sama.com/> (primary)

Quote: "Object Tracking and Temporal Attributes ... Utilize linear interpolation between frames ... player tracking for sports analytics"

Verifier evidence (high): CLAIM CONFIRMED against primary/vendor sources. (Correcting an earlier premature StructuredOutput call I made before web results returned ? web access DID work this

turn.)

VERBATIM PRIMARY-SOURCE CONFIRMATION:

1. "Object Tracking and Temporal Attributes" ? confirmed THREE times via independent WebFetch of sama.com: it is the literal label of the "Video Annotation" service in the site's Data Annotation navigation menu (alongside Image Annotation "Object Detection, Localization and Segmentation"). This is the exact phrase in the claim.
2. "linear interpolation between frames" ? confirmed via Sama-owned pages (sama.com/2d-video-annotation-tool-services, sama.com/blog/object-tracking-in-samahub-with-frame-level-labeling, docs.sama.com/en_US/annotation-products/video-annotation): "Sama offers ML-assisted or linear interpolation between frames to power up object tracking with faster high-quality annotations" and "Sama's built-in automated interpolation between video frames."
3. "static and dynamic annotations" / "complex object changes over time" ? supported: Sama "supports dynamic labeling for object characteristics that change over time, such as visibility percentage of an object or specific characteristics like emotions"; "frame-level labeling technology allows unique objects to be dynamically tracked... changes in position and pose captured with cuboids, polygons and bounding boxes."
4. "player tracking for sports analytics" ? confirmed as a stated Sama use case ("Sama's platform assists in player tracking for sports analytics, object recognition for autonomous vehicles, and medical image analysis"), plus a dedicated Sama blog "Hitting Revenues Out of the Park... Data for Sports Analytics" (sama.com/blog/data-labeling-challenges-for-sports-analytics).

CURRENCY: Sources are current ? sama.com pages reference a "Data Annotation Guide 2026," so this reflects Sama's 2026 offering, not stale info.

SCOPE/HEDGE CHECK: The claim is accurately scoped. It says these features are "directly relevant to badminton rally segmentation" and explicitly hedges "though no detailed sports portfolio is shown." It does NOT falsely assert Sama provides native temporal event-SEGMENT (start/end-timestamp CSV) labeling. WORTH FLAGGING for the decision-maker (does not refute the claim): Sama's confirmed capability is frame-by-frame SPATIAL object/player tracking via bounding-box interpolation, which is a different primitive from temporal-action SEGMENT annotation (marking [start,end) rally ranges). A managed vendor like Sama could almost certainly run rally-segment timestamping as a custom workflow, but the cited features prove tracking, not an out-of-the-box rally-segment product. The claim's own hedge already captures this gap honestly.

VERDICT: Quote-supported, primary-sourced (vendor's own site + docs), current, and not marketing overreach. Source quality matches claim strength.

[7] Sama holds ISO 27001 certification and is GDPR-compliant as a data processor (also CCPA, ISO 9001, ISO 42001, TISAX), but SOC

Vote: 2-1 · Source: <https://www.sama.com/> (primary)

Quote: "ISO 27001 ... GDPR compliant as a data processor ... AICPA SOC-2 certification in progress"

Verifier evidence (high): CORRECTING MY EARLIER ERRONEOUS VOTE: I initially refuted this believing my tools failed, but the tool results actually returned data ? and a direct fetch of the correct primary source CONFIRMS the claim. (The claim's "Source: <https://www.sama.com/>" pointed at the homepage and /security & /trust-center 404; the real trust page is <https://www.sama.com/data-security-and-trust>, last published 2026-05-28, 207KB.)

Element-by-element verification against that primary page (WebFetch quotes):

- ISO 27001 -> listed/certified ("ISO 27001"). CONFIRMED.
- GDPR as data processor -> exact verbatim "GDPR compliant as a data processor" + "EU GDPR COMPLIANT". CONFIRMED.
- CCPA -> "CCPA compliant, to ensure all information shared with us, including PII, is kept secure." CONFIRMED.
- ISO 9001 -> "ISO 9001" certified. CONFIRMED.
- TISAX -> "TISAX". CONFIRMED.
- SOC 2 only IN PROGRESS (not certified) -> page verbatim: "AICPA SOC-2" with notation "soc-2 certification in progress". This is the claim's central assertion and it is DIRECTLY CONFIRMED.
- NO mention of India DPDP Act 2023 -> WebFetch explicitly: "India DPDP / DPDP Act 2023: NOT mentioned on this page." CONFIRMED. (Independent search confirms DPDP is an India-only regime, substantive provisions effective May 2027 ? a global vendor's trust page omitting it is exactly what the claim flags as a gap.)

ONE MINOR INACCURACY (not fatal): The claim lists "ISO 42001" among the certifications Sama

HOLDS, but the primary page actually shows "ISO 42001 in progress" ? i.e. NOT yet certified. So the claim slightly overstates ISO 42001 status. This does not affect the claim's thesis (SOC 2 gap + DPDP gap for an India-focused buyer); it merely understates how much is "in progress." The two load-bearing, decision-relevant assertions (SOC 2 in-progress; DPDP unmentioned) are both fully supported by the vendor's own current trust page, which is the appropriate primary source. Net: claim is substantively accurate, current, and well-sourced -> NOT refuted. (Recommend the synthesis note that ISO 42001 is in-progress, not certified.)

[8] CloudFactory's listed video use cases include action segmentation and human event recognition, the temporal-labeling task class

Vote: 2-1 · Source: <https://www.cloudfactory.com/video-annotation> (primary)

Quote: "Scene understanding and classification; Action segmentation; Human event recognition; Action detection and surveillance"

Verifier evidence (medium): The claim is well-supported by CloudFactory's own (primary) content, and the temporal-labeling characterization is accurate. (1) EXACT-PHRASE MATCH: WebFetch of CloudFactory's own page <https://www.cloudfactory.com/blog/accelerated-annotation-adds-keypoint-annotation-and-video-understanding> explicitly returns all four phrases verbatim ? "Scene understanding and classification", "Action segmentation", "Human event recognition", "Action detection and surveillance" ? under CloudFactory's "video understanding" capability, defined as "identification or classification of activities, actions, or events within streaming video." (2) INDEPENDENT CORROBORATION: a [site:cloudfactory.com](https://www.cloudfactory.com) search returns CloudFactory's Model Monitoring solution page stating they support "human activity recognition, action recognition, and event recognition." (3) TEMPORAL-TASK MAPPING IS CORRECT: "action segmentation" is, by standard computer-vision definition, partitioning a video into contiguous time intervals each labeled with an action (temporal segmentation with start/end boundaries) ? exactly the rally [start,end) task; "human event recognition"/"event recognition" likewise denotes labeling events over time spans. So the claim's assertion that this is "the temporal-labeling task class the badminton golden-set work requires" is substantively right, not an overreach. NOTE/MINOR FLAW: the specifically cited URL <https://www.cloudfactory.com/video-annotation> returned HTTP 404 via WebFetch on multiple attempts, so I could not confirm the quote on that exact page ? but the identical four-phrase taxonomy is verified verbatim on a different CloudFactory-owned page, so the quote is genuine vendor content (at worst the URL is mis-cited, not the substance). This is a primary, current source matching a non-extraordinary capability claim; refutation is not warranted.

[9] iMerit provides video annotation including object tracking, activity recognition, and behavioral analysis, and its Ango Hub present

Vote: 3-0 · Source: <https://imerit.net/services/video-annotation/> (primary)

Quote: "iMerit's Video Annotation solution built on Ango Hub presents users with a timeline view showcasing annotations on the video, with frame-specific classifications allowing annotators to classify individual frames, and the capability to seamlessly add or delete keyframes."

Verifier evidence (high): VERDICT: NOT REFUTED. The claim is well-supported by iMerit's own primary sources, current, and appropriately hedged. (Correcting my erroneous first call, which wrongly assumed network failure; curl returned exit 0 / 65KB and searches returned rich iMerit content.)

(1) Supporting quote is verbatim-accurate. The timeline/frame-specific-classifications/add-or-delete-keyframes language recurs across iMerit-controlled pages and Ango Hub docs.

(2) "Object tracking, activity recognition, and behavioral analysis" is iMerit's OWN wording, not author invention: iMerit video-annotation service copy states "iMerit provides frame-by-frame object tracking, activity recognition, and behavioral analysis through bounding boxes, polygons, and semantic segmentation" (surfaced repeatedly from imerit.net video-annotation-services / data-labeling pages).

(3) The claim's qualifier (video framed primarily as frame-by-frame SPATIAL labeling, not [start,end) timestamp CSVs) is supported by the same "...through bounding boxes, polygons, and semantic segmentation" sentence, the Ango Hub tool list ("Bounding Box, Rotated Bounding Box, Polygon, Polyline, Segmentation, Point and Brush"), and the keyframe-interpolation mechanic. WebFetch of the Ango Hub product page explicitly confirmed: "No mention of [time range marking [start,end) intervals / temporal event segment labeling / CSV timestamp export]... focused on frame-level and frame-range spatial annotations."

(4) The hedge "capabilities adaptable to temporal rally-segment labeling" is conservative and fair. The Ango Hub video-labeling-editor docs (docs.imerit.net) actually show MORE than the claim asserts: frame-specific classifications "behave like a tool-based object with keyframes" and the editor exposes "Set Annotation Start"/"Set Annotation End" buttons, letting annotators

"mark event start/end times independently of spatial tracking" ? i.e. event/segment labeling without a spatial box. So the claim if anything understates capability; it does not overreach.

Source quality (vendor site + official Ango Hub docs) matches the claim's modest descriptive strength. No contradicting source found.

[10] iMerit publicly claims strong security/compliance certifications (SOC 2, ISO 27001, GDPR), which covers most of the buyer's data

Vote: 3-0 · Source: <https://imerit.net/services/video-annotation/> (primary)

Quote: "iMerit is SOC2, ISO 27001, GDPR, HIPAA, and TISAX compliant."

Verifier evidence (high): CORRECTION to my earlier call (which fired before the primary compliance page finished loading and relied on a WebSearch summary that happened to omit SOC 2). The authoritative primary source ? iMerit's own Compliance and Certifications page (<https://imerit.ai/compliance-and-certifications/>, where imerit.net 301-redirects) ? explicitly names ALL of: SOC 2 ("SOC2 attestation is a widely recognized gold standard for data security"), ISO 27001 ("The International Organization for Standardization 27001 Standard (ISO 27001)"), GDPR ("iMerit welcomes GDPR as an important step forward..."), HIPAA, ISO 9001:2015, and TISAX. So the claim's core assertion ? that iMerit publicly claims SOC 2 + ISO 27001 + GDPR (plus HIPAA, TISAX) ? is directly supported by the vendor's own primary source. Multiple independent searches corroborate a SOC 2 Type 2 attestation. The claim's supporting quote ("iMerit is SOC2, ISO 27001, GDPR, HIPAA, and TISAX compliant.") is a faithful paraphrase of the page's content rather than a verbatim sentence, but every certification it lists is substantiated. The second half of the claim is also verified: the same compliance page does NOT mention India's DPDP Act 2023 or a DPA, and every targeted search for "iMerit + DPDP / Digital Personal Data Protection" returned only generic DPDP material with no iMerit connection ? so "DPDP posture and DPA willingness remain unverified" is accurate. Minor caveats (do not rise to refutation): SOC 2 is technically an attestation, not a certification, and the buyer should still request the actual SOC 2 / ISO 27001 reports and a signed DPA rather than relying on website claims. Net: the claim is well-supported by primary sources, current, and appropriately hedged on the DPDP gap.

[11] FlexiBench is the data-annotation platform/managed-workforce arm of Indika AI (its parent company), which is an India-based en

Vote: 3-0 · Source: <https://www.flexibench.io/services/video-annotation> (primary)

Quote: "At Indika (our parent company), we learned early that models are only as good as the data they train on. ... It was founded in 2021 and is based in Mumbai, India."

Verifier evidence (high): NOT REFUTED ? the claim's substantive facts are all well-supported by primary/authoritative sources, though the exact supporting quote is misattributed. (1) FlexiBench = Indika's annotation arm: CONFIRMED ? [flexibench.io](https://www.flexibench.io) home page states "At Indika (our parent company), we learned early that models are only as good as the data they train on. ... [we built] FlexiBench" (WebFetch + curl of [flexibench.io/](https://www.flexibench.io/), LEN 11395). indikaai.com links to [flexibench.io](https://www.flexibench.io) as its annotation platform. (2) India-based: CONFIRMED. (3) HQ Mumbai, NOT Bangalore: CONFIRMED by Indika's OWN /about page ? "Where We Work / Headquartered in India... 01 Mumbai Andheri (West)" (Mumbai listed first), and contact page lists Mumbai first among offices; corroborated by registry CIN U72900MH2021PTC361133 (MH=Maharashtra/Mumbai), inc42 ("Based Out Of Mumbai"), Zaubacorp/TheCompanyCheck registered address "Andheri(W), Mumbai, Maharashtra 400053". (4) Founded 2021: CONFIRMED by CIN ...2021PTC..., YNOS ("founded in 2021 by Hardik Kumar Dave & Urvy Vishal Dave"), Tracxn, multiple sources. CAVEATS: (a) The claim's quoted sentence "It was founded in 2021 and is based in Mumbai, India" does NOT appear on the cited video-annotation page (which now 404s; substring check [mumbai/2021/founded](https://www.flexibench.io/mumbai/2021/founded) all False on the page text that was retrievable) ? the quote is misattributed/likely paraphrased, a sourcing error but not a factual one. (b) Indika AI was acquired by PanScience Innovations (PSI) in Aug 2025 per WebSearch, so "parent company" framing may now be one level removed ? worth flagging but doesn't change the India/Mumbai/2021 conclusion. Net: the operative conclusion (India presence confirmed; HQ Mumbai not Bangalore; FlexiBench = Indika AI under a different brand) is accurate.

[12] FlexiBench explicitly supports temporal/event video annotation ? marking exact start/stop points of events on a video timeline ?

Vote: 1-1 · Source: <https://www.flexibench.io/services/video-annotation> (primary)

Quote: "Activity & Event Annotation: Annotators label distinct events or actions that occur in specific time segments of the video... Time-coded tagging with playback tools support frame-accurate tagging and playback scrubbing to help annotators mark exact start and stop points with minimal drift."

Verifier evidence (high): SUPPORTED, with a citation-hygiene caveat. The claim ("FlexiBench explicitly supports temporal/event video annotation ? marking exact start/stop points of events on a video timeline... not merely static image boxes") is a faithful, non-overreaching restatement of the supporting quote. The load-bearing language is independently corroborated on FlexiBench's own domain: WebSearch surfaced FlexiBench-hosted pages ([flexibench.io/blog/event-](https://www.flexibench.io/blog/event-)

detection-annotating-significant-occurrences and human-activity-recognition-via-video-annotation) reproducing verbatim "Annotators label distinct events or actions that occur in specific time segments of the video" and "Frame-accurate tagging and playback scrubbing to help annotators mark exact start and stop points with minimal drift," plus "multi-layer annotation timelines" and IAA "tracking on event type, timing." A separate search confirmed FlexiBench advertises "sports tracking" as a supported application. This is genuinely temporal/event labeling (start,end on a timeline), and the site separately lists bounding-box/segmentation/keypoint work, so the claim's static-vs-temporal distinction is real. CAVEATS (qualify weight, do not refute): (1) The exact cited URL <https://www.flexibench.io/services/video-annotation> returned HTTP 404 on direct fetch (both with/without www) ? the precise path may have moved; the corroborating language was found on FlexiBench blog/marketing pages rather than that exact services URL, so the citation should be corrected to a live page. (2) This is self-reported marketing copy from a thin/early-stage vendor (parent: Indika AI); no HQ disclosed on its own site, no named sports clients, no third-party case studies/reviews surfaced. The capability is ADVERTISED, not proven ? a buyer must validate via a paid pilot. But the claim is narrowly scoped to what the vendor advertises, so a vendor-owned page is sufficient source quality for this specific assertion.

[13] Labellerr is operated by Tensor Matics, Inc. with a San Francisco HQ and an India back office in Chandigarh, Punjab (and Noida per third-party listings)

Vote: 3-0 · Source: <https://www.labellerr.com/> (primary)

Quote: "HQ: "44, Tehama St, San Francisco, CA USA 94107" ... India Office: "2nd Floor, Innovation Mission Punjab, Kissan Vikas Chamber...Chandigarh, Punjab 140306""

Verifier evidence (high): CLAIM HOLDS (primary sources). (1) Legal entity: [labellerr.com/terms-and-conditions](https://www.labellerr.com/terms-and-conditions) names the operating entity as "Tensor Matics, Inc." (US "Inc."). (2) HQ: [labellerr.com/contact-us](https://www.labellerr.com/contact-us) shows a US/San Francisco, California presence ? consistent with the quoted "44 Tehama St, San Francisco, CA 94107." (3) India back office: the contact page shows an office in the Chandigarh/Mohali tri-city (Punjab) region at "Innovation Mission Punjab," postal range 140xxx ? directly matching the quoted Chandigarh address. (4) The global LinkedIn URL ([linkedin.com/company/labellerr](https://www.linkedin.com/company/labellerr)) redirects to in.linkedin.com, corroborating a real India presence. (5) Tensor Matics Inc has independent third-party profiles (ZoomInfo, Bloomberg), consistent with the "Noida per third-party listings" sub-claim ? which the claim itself appropriately HEDGES as third-party-sourced rather than asserting as primary fact. (6) Most importantly for the research goal: NO Bangalore office appears on any Labellerr page (contact, about-us, terms). The load-bearing conclusion ? "India presence exists but is NOT Bangalore" ? is well-supported. Skeptical checks all pass: quote supports the claim, no contradicting source found, primary-source quality matches claim strength, no staleness or marketing-fluff issues. Minor caveat: I could not independently re-confirm the exact Noida third-party listing text (ZoomInfo/Bloomberg behind anti-bot walls), but the claim correctly flags it as unverified third-party data, so this does not constitute a refutation.

[14] Labellerr is primarily a self-serve video/image annotation SaaS PLATFORM with an optional managed-services add-on ("professional data labeling by experts")

Vote: 2-0 · Source: <https://www.labellerr.com/> (primary)

Quote: "SaaS + Managed Services hybrid. Self-service platform with optional "Annotation Services: Professional data labeling by experts.""

Verifier evidence (high): Claim survives refutation; corroborated by primary source + independent aggregators. (1) SUPPORTED BY SOURCE: Labellerr's homepage (<https://www.labellerr.com/>) positions the product as an AI-powered annotation PLATFORM (e.g., "The Smartest Way to Build Computer Vision Models", auto-labeling/model-assisted tooling) AND separately advertises an "Annotation Services" / expert-annotator option presented alongside the self-serve tool ? exactly the "tool-plus-services hybrid" the claim describes. (2) INDEPENDENT CORROBORATION: G2 categorizes Labellerr under "Data Labeling Software / Training Data Software" (i.e., a platform/tool), and Crunchbase/company descriptions identify it as a SaaS data-annotation platform (product of Tensor Matics, Inc.) that also offers managed/expert annotation as an option. Multiple independent searches converge on the same "self-serve SaaS platform + optional managed annotation services" characterization. (3) NO CONTRADICTION: no credible source positions Labellerr as a pure managed-services/BPO annotation vendor; the "primarily a platform" framing matches G2's software categorization. (4) NOT OUTDATED / NOT a benchmark or marketing-performance claim ? it is a stable, low-extraordinariness product-model description verifiable from the vendor's own site. CAVEATS (do not negate the claim): I could not capture the exact verbatim string "professional data labeling by experts" via a clean page-text fetch (WebFetch returned summarized content), though the substance (platform + expert annotation services) is confirmed across the homepage and several independent sources. Separately, for the broader research goal: Labellerr is India-based but in Mohali/Chandigarh (Punjab), NOT Bangalore ? relevant to vendor selection but irrelevant to this specific product-model claim.

Refuted claims (for transparency)

- "Sama does NOT disclose its headquarters or delivery-center city/country on its public pages and shows no India/Bangalore presence, making it a likely poor fit for the founder's strong local-vendor preference (Sama is widely reported to be San Francisco-HQ'd with East Africa delivery operations ? to be confirmed against secondary sources)." (<https://www.sama.com/>, vote 1-2)
- "CloudFactory explicitly supports TEMPORAL/event video annotation ? including action segmentation, event/activity recognition, and action detection over time ? which is a direct capability match for marking [start,end) rally segments in badminton videos (not merely static frame bounding boxes)." (<https://www.cloudfactory.com/video-annotation>, vote 0-3)
- "Labellerr supports video annotation but its public site does NOT specify temporal/event/segment (start-end timestamp) labeling, action labeling, or sports-video experience for that task, so the exact rally-segmentation capability is UNVERIFIED despite a generic "Sports Vision" solution link." (<https://www.labellerr.com/>, vote 0-0)
- "TELUS International acquired Playment, a Bangalore/Bengaluru-based data-annotation company, announced on July 6, 2021 ? confirming the acquisition referenced in the research question's vendor list (Playment is no longer an independent vendor)." (<https://www.telusdigital.com/about/news-and-events/news/2021/playment-acquisition>, vote 0-0)
- "Playment's core capability covered VIDEO annotation (not just static images) including LiDAR/3D and sensor data, and its standalone product offered object tracking for video ? directly relevant to the temporal/video-segment annotation need, though its documented video tooling is frame-level object tracking rather than rally start/end timestamp segmentation." (<https://www.telusdigital.com/about/news-and-events/news/2021/playment-acquisition>, vote 0-0)
- "TELUS Digital provides image/video annotation including object detection, segmentation, landmark recognition, and movement prediction, confirming video annotation is a stated capability, though the page does not explicitly mention temporal/event timestamp segmentation of video." (<https://www.telusdigital.com/solutions/data-for-ai-training/data-annotation-services>, vote 0-0)
- "TELUS Digital states its facilities are ISO 27001 certified and the company is SOC 2 compliant and TISAX certified, establishing a documented security/compliance posture, though GDPR and India's DPDP Act 2023 are not mentioned on this page." (<https://www.telusdigital.com/solutions/data-for-ai-training/data-annotation-services>, vote 0-0)
- "iMerit holds SOC 2 Type 2 attestation and undergoes regular audits, indicating an externally-attested data security program (relevant to the security/DPA vendor-screening criterion)." (<https://imerit.net/compliance-and-certifications/>, vote 0-0)
- "iMerit is ISO 27001 certified, an information security standard that covers its office sites, development centers, support centers and data centers." (<https://imerit.net/compliance-and-certifications/>, vote 0-0)
- "iMerit's security policies and processes have been evaluated for GDPR compliance, addressing one of the required data-protection frameworks for the golden-set vendor." (<https://imerit.net/compliance-and-certifications/>, vote 0-0)

Instructions

1. Identify claims that say the same thing ? merge them, combine their sources.
2. Group related claims into coherent findings. Each finding should directly address the research question.
3. Assign confidence per finding: high (multiple primary sources, unanimous votes), medium (secondary sources or split votes), low (single source or blog-quality).
4. Write a 3-5 sentence executive summary answering the research question.
5. Note caveats: what's uncertain, what sources were weak, what time-sensitivity applies.
6. List 2-4 open questions that emerged but weren't answered.

Structured output only.

2. assistant (2026-05-30T09:41:07.300Z)

You've hit your session limit · resets 6:30pm (Asia/Calcutta)