

claude-windows / aa329f09-28b7-4acc-b062-98ec7e905abc

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\aa329f09-28b7-4acc-b062-98ec7e905abc\subagents\workflows\wf_ddd29371-8c2\agent-afe7e4d9ebc8a3a19.jsonl`  
- SHA-256: `5d3ca0ff64f86ccb6b384fa9e576bc4d031d6d450ddc60c351fa4f019926d80e`  
- Source modified: `2026-05-30T03:12:45+00:00`  
- Imported at: `2026-07-05T16:48:25+00:00`  
- project: `wf_ddd29371-8c2`  
- session_id: `aa329f09-28b7-4acc-b062-98ec7e905abc`
```

Transcript

1. user (2026-05-30T03:12:08.725Z)

Adversarial Claim Verifier (voter 3/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

RESEARCH GOAL: Identify and eliminate every dependency that could legally or economically hinder monetizing the "badminton-highlight-indexer" project as a HOSTED API SaaS. Business model: the project runs server-side on a cloud GPU VM (NVIDIA); users upload their own badminton videos and receive auto-generated highlight reels. CRITICAL CONSTRAINT: we do NOT distribute the code or any binary to users ? everything stays server-side ? so software that is merely *hosted* on the server (invoked as a separate process, e.g. ffmpeg) is acceptable where distribution-triggered copyleft (GPL/LGPL) would NOT be triggered, BUT network/SaaS copyleft (AGPL §13) and external metered-API terms STILL apply. Preference order: permissively-licensed (Apache-2.0/MIT/BSD) and royalty-free, self-hostable components. Paid commercial licenses (e.g. Ultralytics Enterprise, codec patent licenses) are acceptable ONLY as clearly-flagged fallbacks. Output must be decision-grade with licenses verified against PRIMARY sources (actual GitHub LICENSE files, official pricing/terms pages), because this drives a real business decision ? flag any claim you could not verify.

CURRENT STACK TO AUDIT (verify each license + its SaaS implication):

- ultralytics (Python lib) + the bundled YOLOv8n.pt model weights ? believed AGPL-3.0. Confirm the license of BOTH the code and the pretrained weights, whether the AGPL network clause is triggered by a closed-source hosted service that uses ultralytics server-side, and the existence/cost/terms of the Ultralytics Enterprise license.
- google-genai SDK + the Google Gemini API (model gemini-2.5-flash). Distinguish the SDK license (Apache?) from the SERVICE terms. Confirm: (a) is commercial use of outputs allowed in a paid product; (b) does Google train on / human-review prompt+video data on the PAID/billed tier vs the free tier; (c) abuse-monitoring / data-retention terms; (d) current per-token / per-video pricing for gemini-2.5-flash (input video tokens-per-second, output) so we can sketch per-video cost.
- ffmpeg / ffprobe (system binaries, invoked as separate processes). Clarify GPL vs LGPL builds and whether pure server-side SaaS use (no distribution to users) triggers any source-disclosure obligation. SEPARATELY and importantly: video CODEC PATENT royalties for a commercial service that DECODES user-uploaded H.264/AVC and H.265/HEVC (e.g. GoPro footage) and ENCODES output ? covering MPEG-LA / Via-LA (AVC), Access Advance + Via-LA + others (HEVC) pools, AAC audio (Via-LA), what activities trigger royalties, any free thresholds (e.g. the AVC sub-100k or free-internet-broadcast provisions), and whether a small SaaS realistically incurs these. Then royalty-FREE codec alternatives for the OUTPUT we generate: AV1, VP9, Opus ? and their encoder maturity/speed on a GPU VM.
- torch, torchvision ? confirm BSD; flag that some torchvision *pretrained weights* may carry separate licenses.
- opencv-python ? confirm Apache-2.0 and that the PyPI wheel's bundled components are clean for commercial server use.

- lapx (a fork of "lap" for ByteTrack association) ? confirm license.
- Quick permissive confirmation for: fastapi, uvicorn, pydantic, jinja2, python-dotenv, numpy.

PERMISSIVE ALTERNATIVES TO RESEARCH (all must be GPU-VM-friendly and usable in a closed-source commercial SaaS):

1. Object detection to replace ultralytics/YOLOv8 for player/person detection (and possibly shuttle): compare YOLOX (Apache-2.0), RT-DETR / RT-DETRv2 (the Baidu/Apache implementations ? NOT the Ultralytics port), PP-YOLOE, MMDetection (Apache-2.0), Detectron2 (Apache-2.0), torchvision built-in detectors (Faster R-CNN/RetinaNet/FCOS/SSD, BSD), RF-DETR (Roboflow), and D-FINE/DEIM. For EACH, verify BOTH the code license AND the pretrained-weights license (some "open" detectors like YOLO-NAS have restrictive weight licenses ? confirm). Note accuracy/speed trade-offs for person detection on sports video.
2. Shuttle/ball trajectory tracking ? VERIFY the actual repo licenses (commercial-use viability) of: WASB / WASB-SBDT ("Widely Applicable Strong Baseline" for sports ball detection), TrackNetV2 and TrackNetV3, MonoTrack, and any permissively-licensed alternative for tiny-fast-object tracking. This is the planned shuttle-tracking substrate, so its commercial-use license is high-stakes.
3. LLM / semantic rally-boundary step ? research ALL THREE strategies (the user wants all three compared):
 - (a) ELIMINATE the hosted LLM entirely: permissively-licensed temporal action segmentation / video models suitable for learning rally boundaries offline (e.g. MS-TCN/MS-TCN++, ASFormer, ActionFormer, and similar) ? verify licenses and whether they fit a from-annotations learned segmenter.
 - (b) SELF-HOST an open-weights vision-language model that can ingest video frames for semantic boundaries: compare Qwen2.5-VL, InternVL2/2.5, LLaVA-OneVision/LLaVA-Video, MiniCPM-V, VideoLLaMA ? for EACH verify the WEIGHTS license for commercial use (Apache-2.0 vs custom "research-only"/"non-commercial"/community licenses with MAU caps) and the GPU VRAM needed per size tier.
 - (c) KEEP Gemini but de-risk: summarize commercial terms, data-usage on paid tier, rate limits, and per-video cost at scale, plus other hosted-API alternatives (OpenAI, Anthropic) and how their commercial/data terms compare.

DELIVERABLE: A cited report containing (1) a per-dependency RISK TABLE ? columns: dependency, license, does-SaaS-hosting-trigger-obligations?, monetization verdict (CLEAR / WATCH / BLOCKER), recommended action; (2) a recommended FULLY-PERMISSIVE target stack that eliminates all BLOCKERS; (3) paid-license fallbacks where relevant; (4) a rough per-video processing-cost sketch on a GPU VM for each of the three LLM strategies; (5) an explicit list of the codec/patent considerations and the royalty-free output-codec recommendation. Prioritize accuracy and primary-source citations over breadth.

Claim under review

"On Paid Services (including the paid quota of the Gemini API), Google does not use prompts (including system instructions, cached content, and files such as images, videos, or documents) or responses to improve its products. This is the key data-usage assurance for a commercial SaaS billed via the paid Gemini API."

****Source:**** <https://ai.google.dev/gemini-api/docs/data-usage> (primary)

****Supporting quote:**** "When you use Paid Services, including, for example, the paid quota of the Gemini API, Google doesn't use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve our products"

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-05-30T03:12:11.871Z)

I'll verify this claim about Google's data-usage terms for Paid Services. Let me check primary sources.

3. user (2026-05-30T03:12:29.897Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

4. user (2026-05-30T03:12:33.477Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

5. user (2026-05-30T03:12:35.103Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

6. user (2026-05-30T03:12:38.572Z)

Web search results for query: "Gemini API paid tier data usage Google does not use prompts improve products exceptions abuse monitoring"

Links: [{"title":"Billing | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/billing"}, {"title":"Zero data retention in the Gemini Developer API | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/zdr"}, {"title":"How Gemini for Google Cloud uses your data | Google Cloud Documentation","url":"https://docs.cloud.google.com/gemini/docs/discover/data-governance"}, {"title":"Abuse monitoring | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/usage-policies"}, {"title":"Is Gemini API Free? (2025) - Spurn","url":"https://www.spurnow.com/en/blogs/is-gemini-api-free"}, {"title":"Gemini API Additional Terms of Service | Google AI for Developers","url":"https://ai.google.dev/gemini-api/terms_preview"}, {"title":"Service update: mitigating abuse and prioritizing traffic · google-gemini/gemini-cli · Discussion #22970","url":"https://github.com/google-gemini/gemini-cli/discussions/22970"}, {"title":"Clarify use of user's code for training within the free quota on paid service · Issue #1472 · google-gemini/gemini-cli","url":"https://github.com/google-gemini/gemini-cli/issues/1472"}, {"title":"Data Logging and Sharing | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/logs-policy"}, {"title":"Gemini API Free Tier 2026: 1,500 Req/Day, 1M TPM ? No Card - TokenMix Blog","url":"https://tokenmix.ai/blog/gemini-api-free-tier-limits"}]

Here are the search results for your query about Gemini API paid tier data usage, Google's use of prompts, and abuse monitoring:

Key Findings:

Data Usage for Paid Services:

When you use Paid Services, Google does not use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve our products.

****Exceptions for Abuse Monitoring:****

For Paid Services, Google logs prompts and responses for a limited period of time solely for detecting violations of the Prohibited Use Policy. More specifically, Google retains the following data for fifty-five (55) days for the purposes of detecting and preventing violations of the Prohibited Use Policy to maintain the safety and security of the Services: Prompts, Contextual Information, and Output.

****Important Note on Abuse Data Usage:****

When data is logged for abuse monitoring, it is used solely for the purpose of policy enforcement and preventing policy violations. It is not used to train or fine-tune any AI/ML models besides those used specifically for policy enforcement.

****Grounding Features Exception:****

When using Grounding with Google Search, Google stores prompts, contextual information, and generated output for thirty (30) days for the purposes of creating grounded results and search suggestions. This stored information may be used for debugging and testing of systems that support grounding.

****Zero Data Retention Option:****

When your request for ZDR for a particular project is approved, all user content (prompts and responses) and identifiable metadata (such as IP addresses and Google Account IDs) are cleared prior to logging.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

7. user (2026-05-30T03:12:41.399Z)

The server returned HTTP 404 Not Found.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

8. user (2026-05-30T03:12:43.506Z)

Google's Data Usage Policies: Paid vs. Unpaid Services

Unpaid Services

Google ****does**** use your data to improve products:

> "Google uses the content you submit to the Services and any generated responses to provide, improve, and develop Google products"

Human review occurs:

> "human reviewers may read, annotate, and process your API input and output"

However, Google takes privacy steps:

> "This includes disconnecting this data from your Google Account, API key, and Cloud project before reviewers see it"

Paid Services

Google ****does not**** use prompts/responses for improvement:

> "Google doesn't use your prompts...or responses to improve our products"

Data retention is limited:

> "Google logs prompts and responses for a limited period of time, solely for detecting and

preventing violations"

The stated purposes are narrow:

> "solely for detecting and preventing violations of the Prohibited Use Policy to maintain the safety and security"

Key Exception

For European Economic Area, Switzerland, or UK users, the restrictive Paid Services data terms apply **even to free offerings**.

9. user (2026-05-30T03:12:45.487Z)

Web search results for query: "'Gemini API' paid services 'doesn't use your prompts' data usage terms"

Links: [{"title": "Gemini API Additional Terms of Service | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/terms"}, {"title": "Billing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/billing"}, {"title": "Is Gemini API Free? (2025) - Spurn", "url": "https://www.spurnow.com/en/blogs/is-gemini-api-free"}, {"title": "Gemini CLI: Terms of Service and Privacy Notice | gemini-cli", "url": "https://google-gemini.github.io/gemini-cli/docs/tos-privacy.html"}, {"title": "Gemini Business Terms of Service | Google Cloud", "url": "https://cloud.google.com/terms/gemini-enterprise/business"}, {"title": "Gemini CLI: License, Terms of Service, and Privacy Notices | Gemini CLI", "url": "https://geminicli.com/docs/resources/tos-privacy/"}, {"title": "Gemini CLI: Quotas and pricing | Gemini CLI", "url": "https://geminicli.com/docs/resources/quota-and-pricing/"}, {"title": "Gemini API & Data Privacy: What Google's AI Terms Mean for You in 2025", "url": "https://redact.dev/blog/gemini-api-terms-2025"}, {"title": "Abuse monitoring | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/usage-policies"}, {"title": "Gemini Developer API pricing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/pricing"}]

Based on the search results, here's what Google's data usage terms say about Gemini API paid services:

Data Usage for Paid Services

When you use Paid Services of the Gemini API, Google doesn't use your prompts (including associated system instructions, cached content, and files such as images, videos, or documents) or responses to improve our products, and will process your prompts and responses in accordance with the Data Processing Addendum for Products Where Google is a Data Processor.

Limited Logging for Abuse Detection

However, data isn't completely excluded from all logging. For Paid Services, Google logs prompts and responses for a limited period of time, solely for detecting and preventing violations of the Prohibited Use Policy to maintain the safety and security of the Services, and any required legal or regulatory disclosures. Additionally, Google retains the following data for fifty-five (55) days for the purposes of detecting and preventing violations of the Prohibited Use Policy to maintain the safety and security of the Services, and any required legal or regulatory disclosures: Prompts: The text prompts you submit to the API.

When data is logged for abuse monitoring, it is used solely for the purpose of policy enforcement and preventing policy violations. It is not used to train or fine-tune any AI/ML models besides those used specifically for policy enforcement.

Key Difference from Free Tier

In contrast, on the free tier, prompts and responses can be used (with human review) to improve Google's products. On paid tiers, prompts and responses are not used to improve Google's models and are processed under Google's data processing addendum.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.