

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_076beba8-59c\agent-a4b2c50f1644d7d0b.jsonl`  
- SHA-256: `c562f28803a44a5b920ee230f09b5ebd9106cf0051797cefc5bdda78bc365592`  
- Source modified: `2026-05-31T04:22:55+00:00`  
- Imported at: `2026-07-05T16:48:20+00:00`  
- project: `wf_076beba8-59c`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T04:22:03.423Z)

Adversarial Claim Verifier (voter 1/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Identify the best LOCAL (self-hosted, single NVIDIA CUDA GPU) vision / document-AI / OCR models ? and, as a clearly-flagged approval-gated fallback only, paid cloud API services ? for parsing BANK STATEMENT PDFs into structured transaction tables (date, narration/description, withdrawal/debit, deposit/credit, running balance), for a LOCAL-FIRST personal-finance tool ("muneem").

Target documents: Indian bank statements (HDFC, Axis, ICICI, AU Small Finance) ? dense, multi-column, often SEMI-RULED or borderless tables, sometimes multiple accounts per statement; both born-digital (text layer present) and the harder cases where pdfplumber's extract_tables/coordinate clustering is unreliable (e.g. Axis, whose column layout varies between statements).

HARD CONSTRAINTS (a model that violates these is unusable for us):

1. MUST run 100% LOCALLY for statement contents (highly sensitive). NO cloud APIs by default. Paid/cloud is acceptable ONLY as an explicitly-approved fallback ? still report the best cloud options but label them clearly as "cloud/sensitive ? approval-gated, not default".
2. MUST be usable from Python (transformers / vLLM / Ollama / ONNX Runtime / native lib).
3. License MUST be permissive & commercially safe for BOTH the CODE and the MODEL WEIGHTS: MIT, BSD, Apache-2.0, ISC, HPND. AVOID and explicitly FLAG anything GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, or custom "community"/acceptable-use licenses with restrictions. This is critical ? many document-AI models and their training datasets are non-commercial (e.g. LayoutLMv3 weights are CC-BY-NC; Surya and Marker use GPL-or-commercial dual licensing; several VLM weights ship custom restrictive licenses). Verify the WEIGHTS license, not just the repo license.
4. Hardware: a single consumer/prosumer NVIDIA GPU. Give VRAM tiers (~8-12 GB consumer, ~16-24 GB) and which models/quantizations fit each.

Evaluate and compare these LOCAL approaches ? for each give: code license, MODEL-WEIGHTS license (state explicitly), latest release / maintenance status as of 2025-2026, VRAM need + quantization options, Python integration path, and reliability specifically on DENSE FINANCIAL TABLES / multi-column statements:

- Open vision-language models (VLMs) for document/table understanding: Qwen2-VL and Qwen2.5-VL (note per-size license ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2 / 2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision (flag the Llama Community License restrictions), Moondream2, GOT-OCR2.0, olmOCR and other OCR-specialized VLMs.
- OCR-specialist / document-AI pipelines: PaddleOCR + PP-Structure (table recognition), docTR, EasyOCR, Tesseract (already used), Microsoft Table Transformer / TATR (check code AND PubTables-1M dataset/weights license), Donut, Surya (FLAG license), Marker (FLAG license), MinerU / PDF-Extract-Kit, RapidOCR.

- Table-structure recognition specifically (cell/grid reconstruction for borderless tables): TATR, PP-Structure, UniTable, etc.

Also assess: for SEMI-RULED / borderless bank tables, is a prompt-driven VLM (page image -> JSON transactions) more reliable than a classic OCR + table-structure pipeline? What are the accuracy/hallucination risks of VLMs on financial NUMBERS (transposed digits, dropped/duplicated rows), and how should output be VERIFIED? (muneem already machine-checks parses by balance reconciliation ? opening + sum(deposits) - sum(withdrawals) == closing, plus per-row running-balance walks and cross-checks against the bank's printed summary ? so any model's output can be automatically validated, which should inform the recommendation.)

Finally, briefly cover the best PAID/CLOUD fallbacks (clearly flagged approval-gated for sensitive data): AWS Textract (AnalyzeExpense / Tables), Google Document AI, Azure AI Document Intelligence, Reducto, LlamaParse, Anthropic/OpenAI vision ? note any self-hosting option and data-handling terms.

Deliver a RANKED shortlist of the most promising LOCAL options for muneem (NVIDIA GPU, Python, commercial-license-safe), each with license (code + weights), VRAM tier, maturity, strengths/weaknesses for dense bank-statement tables, and a concrete recommendation: which local model(s) to prototype for muneem's statement parsing and how they'd combine with the existing deterministic parsers + balance-reconciliation verification (e.g. VLM as a fallback/cross-check when deterministic parsing fails to reconcile). Cite sources and call out any license uncertainty explicitly.

Claim under review

"The Apache-2.0 attribution in the InternVL2.5-4B model card for its Qwen2.5-3B-Instruct component is disputed as incorrect: the actual license of Qwen2.5-3B-Instruct is the Qwen RESEARCH LICENSE, not Apache-2.0."

****Source:**** <https://github.com/OpenGVLab/InternVL/issues/900> (primary)

****Supporting quote:**** "However, the license for Qwen2.5-3B-Instruct is not Apache 2.0, but rather the Qwen RESEARCH LICENSE."

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-05-31T04:22:06.286Z)

I'll verify this claim about the InternVL2.5-4B license attribution dispute. Let me start by examining the primary source and searching for corroborating/contradicting evidence.

3. assistant (2026-05-31T04:22:08.787Z)

I'll fetch the primary source first, then search for corroborating/contradicting evidence in parallel.

4. user (2026-05-31T04:22:15.530Z)

GitHub Issue #900: InternVL2.5-4B License Clarification

****Issue Title:**** "About the license of internVL2.5-4B model"

****Status:**** Closed

****Issue Author:**** NaokiAsatani (opened February 12, 2025)

License Discrepancy

The issue author raised a concern about conflicting license information. According to the issue:

****From HuggingFace page (as cited by author):****

> "This project uses the pre-trained Qwen2.5-3B-Instruct as a component, which is licensed under the Apache License 2.0."

****Author's Correction:****

The author states that Qwen2.5-3B-Instruct is actually licensed under the ****Qwen RESEARCH LICENSE****, not Apache 2.0. They note this research license prohibits commercial use.

Key Question

The author explicitly asked: "I would like to confirm whether the InternVL2_5-4B model is eligible for commercial use under these circumstances."

Response Status

Based on the provided content, ****no response from the InternVL maintainers is shown****. The issue appears closed, but no acknowledgment or clarification from the team is included in this excerpt.

5. user (2026-05-31T04:22:16.149Z)

Web search results for query: "Qwen2.5-3B-Instruct license Apache 2.0 or Qwen Research License"

Links: [{"title": "Qwen2.5-Coder Series: Powerful, Diverse, Practical - Alibaba Cloud Community", "url": "https://www.alibabacloud.com/blog/qwen2-5-coder-series-powerful-diverse-practical_601765"}, {"title": "LICENSE · Qwen/Qwen2.5-32B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-32B-Instruct/blob/main/LICENSE"}, {"title": "Qwen2.5: A Party of Foundation Models! | Qwen", "url": "https://qwenlm.github.io/blog/qwen2.5/"}, {"title": "Qwen API and Models | OpenRouter", "url": "https://openrouter.ai/qwen"}, {"title": "Qwen/Qwen2.5-3B-Instruct · Why a proprietary license?", "url": "https://huggingface.co/Qwen/Qwen2.5-3B-Instruct/discussions/1"}, {"title": "LICENSE · Qwen/Qwen2.5-0.5B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-0.5B-Instruct/blob/main/LICENSE"}, {"title": "LICENSE · Qwen/Qwen2.5-14B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-14B-Instruct/blob/main/LICENSE"}, {"title": "Qwen/Qwen2.5-0.5B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-0.5B-Instruct/tree/main"}]

Based on the search results, I found the answer to your query:

The Qwen2.5-3B-Instruct model is licensed under the Qwen Research License, not Apache 2.0. All Qwen2.5 open-source models, except for the 3B and 72B variants, are licensed under Apache 2.0.

To clarify further: Among the Qwen2.5-Coder series, the 0.5B, 1.5B, 7B, 14B, and 32B models are licensed under Apache 2.0, while the 3B model is under the Qwen Research License.

So the Qwen2.5-3B-Instruct uses the proprietary Qwen Research License rather than Apache 2.0.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. assistant (2026-05-31T04:22:21.285Z)

The primary source confirms the claim's structure, and the WebSearch strongly corroborates it.

Let me verify the actual license file on the Qwen2.5-3B-Instruct HuggingFace repo directly (the authoritative source for the weights license), and also confirm the official Qwen blog's statement about which sizes are non-Apache.

7. user (2026-05-31T04:22:24.574Z)

Model License

According to the page, the license shown for this model is **Qwen Research License** (displayed as "License: qwen-research" in the metadata section near the top of the page).

It is **not** Apache 2.0.

8. user (2026-05-31T04:22:26.462Z)

Qwen2.5 Licensing

According to the blog post, the licensing situation is:

Apache 2.0: All open-source models except the 3B and 72B variants use this license.

3B and 72B models: These are explicitly excluded from Apache 2.0 licensing.

Here's the exact statement:

```
> "All our open-source models, except for the 3B and 72B variants, are licensed under Apache 2.0."
```

The post does not specify what alternative license applies to the 3B and 72B models?only that they are not covered by Apache 2.0.

9. assistant (2026-05-31T04:22:37.853Z)

I now have multiple independent confirmations. Let me run through the verification checklist.

1. Is the claim supported by the quote? Yes. The issue (#900, "About the license of internVL2.5-4B model", opened Feb 12 2025 by NaokiAsatani) explicitly states the HuggingFace card claims Qwen2.5-3B-Instruct is "licensed under the Apache License 2.0," and the author corrects that it is actually the Qwen RESEARCH LICENSE. The supporting quote matches the issue content exactly.

2. Contradicting evidence? None found ? the opposite. The authoritative HuggingFace repo for Qwen2.5-3B-Instruct shows `License: qwen-research`, NOT Apache 2.0. The discussion thread "Why a proprietary license?" corroborates.

3. Source quality for claim strength? Strong. Three independent sources: (a) the primary GitHub issue, (b) the official Qwen2.5 blog stating "All our open-source models, except for the 3B and 72B variants, are licensed under Apache 2.0," and (c) the model card itself showing the qwen-research tag.

4. Outdated? No. Issue is from Feb 2025; the qwen-research license on the 3B weights is still current as of my checks.

5. Marketing/speculation? No ? it's a license-file fact verifiable on the official HuggingFace repo.

One nuance worth flagging on the claim's exact wording: the claim says the Apache-2.0 attribution "is disputed as incorrect." It is more than disputed ? it is **verifiably incorrect**. The InternVL2.5-4B card's statement that the embedded Qwen2.5-3B-Instruct is Apache-2.0 is factually wrong; the 3B weights carry the Qwen Research License (non-commercial). So the claim is not just supported, it understates the case (calling a verified factual error a "dispute"). This makes the claim's core assertion well-founded. The only minor caveat ? the InternVL team did not visibly acknowledge/confirm the error in the thread ? does not weaken the

substantive license fact, which is independently verifiable.

The claim survives. I cannot refute it.

10. user (2026-05-31T04:22:49.037Z)

Structured output provided successfully

11. assistant (2026-05-31T04:22:54.980Z)

I've completed the verification. The claim ****survives**** (refuted=false).

The claim that InternVL2.5-4B's model card incorrectly attributes Apache-2.0 to its embedded Qwen2.5-3B-Instruct component ? when the actual license is the Qwen Research License ? is well-supported by three independent sources:

1. ****Primary**** (GitHub issue #900): The issue author cites the card's Apache-2.0 claim and corrects it to Qwen Research License.
2. ****Authoritative weights source**** (HuggingFace Qwen2.5-3B-Instruct): Shows `License: qwen-research`, not Apache-2.0.
3. ****Official Qwen blog****: "All our open-source models, except for the 3B and 72B variants, are licensed under Apache 2.0."

If anything, the claim **understates** the evidence ? the attribution isn't merely "disputed," it's verifiably incorrect. The one caveat (no visible InternVL-maintainer acknowledgment in the thread) doesn't weaken the underlying license fact, which is independently verifiable.

This is directly relevant to muneem's commercial-license-safety constraint: ****InternVL2.5-4B's weights are tainted by the non-commercial Qwen Research License**** of its 3B base, making them commercially unsafe despite the card's erroneous Apache-2.0 label.