

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_076beba8-59c\agent-a5ce380c9655288e4.jsonl`  
- SHA-256: `14c72e193f7bd80578a86a5445ece15a9283e089708e31e70bd099f1bce3e83d`  
- Source modified: `2026-05-31T04:22:52+00:00`  
- Imported at: `2026-07-05T16:48:20+00:00`  
- project: `wf_076beba8-59c`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T04:22:05.736Z)

Adversarial Claim Verifier (voter 3/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Identify the best LOCAL (self-hosted, single NVIDIA CUDA GPU) vision / document-AI / OCR models ? and, as a clearly-flagged approval-gated fallback only, paid cloud API services ? for parsing BANK STATEMENT PDFs into structured transaction tables (date, narration/description, withdrawal/debit, deposit/credit, running balance), for a LOCAL-FIRST personal-finance tool ("muneem").

Target documents: Indian bank statements (HDFC, Axis, ICICI, AU Small Finance) ? dense, multi-column, often SEMI-RULED or borderless tables, sometimes multiple accounts per statement; both born-digital (text layer present) and the harder cases where pdfplumber's extract_tables/coordinate clustering is unreliable (e.g. Axis, whose column layout varies between statements).

HARD CONSTRAINTS (a model that violates these is unusable for us):

1. MUST run 100% LOCALLY for statement contents (highly sensitive). NO cloud APIs by default. Paid/cloud is acceptable ONLY as an explicitly-approved fallback ? still report the best cloud options but label them clearly as "cloud/sensitive ? approval-gated, not default".
2. MUST be usable from Python (transformers / vLLM / Ollama / ONNX Runtime / native lib).
3. License MUST be permissive & commercially safe for BOTH the CODE and the MODEL WEIGHTS: MIT, BSD, Apache-2.0, ISC, HPND. AVOID and explicitly FLAG anything GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, or custom "community"/acceptable-use licenses with restrictions. This is critical ? many document-AI models and their training datasets are non-commercial (e.g. LayoutLMv3 weights are CC-BY-NC; Surya and Marker use GPL-or-commercial dual licensing; several VLM weights ship custom restrictive licenses). Verify the WEIGHTS license, not just the repo license.
4. Hardware: a single consumer/prosumer NVIDIA GPU. Give VRAM tiers (~8-12 GB consumer, ~16-24 GB) and which models/quantizations fit each.

Evaluate and compare these LOCAL approaches ? for each give: code license, MODEL-WEIGHTS license (state explicitly), latest release / maintenance status as of 2025-2026, VRAM need + quantization options, Python integration path, and reliability specifically on DENSE FINANCIAL TABLES / multi-column statements:

- Open vision-language models (VLMs) for document/table understanding: Qwen2-VL and Qwen2.5-VL (note per-size license ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2 / 2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision (flag the Llama Community License restrictions), Moondream2, GOT-OCR2.0, olmOCR and other OCR-specialized VLMs.
- OCR-specialist / document-AI pipelines: PaddleOCR + PP-Structure (table recognition), docTR, EasyOCR, Tesseract (already used), Microsoft Table Transformer / TATR (check code AND PubTables-1M dataset/weights license), Donut, Surya (FLAG license), Marker (FLAG license), MinerU / PDF-Extract-Kit, RapidOCR.

- Table-structure recognition specifically (cell/grid reconstruction for borderless tables): TATR, PP-Structure, UniTable, etc.

Also assess: for SEMI-RULED / borderless bank tables, is a prompt-driven VLM (page image -> JSON transactions) more reliable than a classic OCR + table-structure pipeline? What are the accuracy/hallucination risks of VLMs on financial NUMBERS (transposed digits, dropped/duplicated rows), and how should output be VERIFIED? (muneem already machine-checks parses by balance reconciliation ? opening + sum(deposits) - sum(withdrawals) == closing, plus per-row running-balance walks and cross-checks against the bank's printed summary ? so any model's output can be automatically validated, which should inform the recommendation.)

Finally, briefly cover the best PAID/CLOUD fallbacks (clearly flagged approval-gated for sensitive data): AWS Textract (AnalyzeExpense / Tables), Google Document AI, Azure AI Document Intelligence, Reducto, LlamaParse, Anthropic/OpenAI vision ? note any self-hosting option and data-handling terms.

Deliver a RANKED shortlist of the most promising LOCAL options for muneem (NVIDIA GPU, Python, commercial-license-safe), each with license (code + weights), VRAM tier, maturity, strengths/weaknesses for dense bank-statement tables, and a concrete recommendation: which local model(s) to prototype for muneem's statement parsing and how they'd combine with the existing deterministic parsers + balance-reconciliation verification (e.g. VLM as a fallback/cross-check when deterministic parsing fails to reconcile). Cite sources and call out any license uncertainty explicitly.

Claim under review

"The Apache-2.0 attribution in the InternVL2.5-4B model card for its Qwen2.5-3B-Instruct component is disputed as incorrect: the actual license of Qwen2.5-3B-Instruct is the Qwen RESEARCH LICENSE, not Apache-2.0."

****Source:**** <https://github.com/OpenGVLab/InternVL/issues/900> (primary)

****Supporting quote:**** "However, the license for Qwen2.5-3B-Instruct is not Apache 2.0, but rather the Qwen RESEARCH LICENSE."

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-05-31T04:22:09.257Z)

I'll verify this claim about the InternVL2.5-4B model card's Apache-2.0 attribution for Qwen2.5-3B-Instruct being disputed. Let me start by loading the web tools.

3. assistant (2026-05-31T04:22:12.806Z)

I'll fetch the primary source first - the GitHub issue.

4. user (2026-05-31T04:22:19.293Z)

Issue Summary

****Core Question:**** Yes, this issue directly addresses a licensing discrepancy regarding Qwen2.5-3B-Instruct used in InternVL2.5-4B.

****The Claim:****

The issue author (NaokiAsatani) states that the Hugging Face model card contains incorrect information. They assert:

> "the license for Qwen2.5-3B-Instruct is not Apache 2.0, but rather the Qwen RESEARCH LICENSE"

The author notes this is problematic because the Qwen Research License restricts commercial use, creating uncertainty about whether InternVL2.5-4B itself is commercially usable.

****Maintainer Response:****

The provided content shows no maintainer confirmation or refutation of this claim. The issue appears to be an open question awaiting clarification from the OpenGVLab team.

****Key Concern:****

There is a documented mismatch between what the InternVL2.5-4B model card claims (Apache 2.0) and what the issue author contends is the actual license of the component (Qwen Research License with commercial restrictions).

5. user (2026-05-31T04:22:19.955Z)

Web search results for query: "Qwen2.5-3B-Instruct license Qwen Research License not Apache 2.0"

Links: [{"title": "LICENSE · Qwen/Qwen2.5-VL-3B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-3B-Instruct/blob/main/LICENSE"}, {"title": "LICENSE · Qwen/Qwen2.5-3B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-3B-Instruct/blob/main/LICENSE"}, {"title": "Qwen2.5 VL 3B Instruct - API Pricing & Providers | OpenRouter", "url": "https://openrouter.ai/qwen/qwen2.5-vl-3b-instruct"}, {"title": "LICENSE · Qwen/Qwen2.5-VL-3B-Instruct at acbce9a6c80e81356680a632166380c8cc1d27ec", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-3B-Instruct/blob/acbce9a6c80e81356680a632166380c8cc1d27ec/LICENSE"}, {"title": "update README & LICENSE · Qwen/Qwen2.5-3B-Instruct-GPTQ-Int4 at 5cabe8e", "url": "https://huggingface.co/Qwen/Qwen2.5-3B-Instruct-GPTQ-Int4/commit/5cabe8e5eacf08566cf6f18a43d1e52547e469d5"}, {"title": "LICENSE · Qwen/Qwen2.5-3B at main", "url": "https://huggingface.co/Qwen/Qwen2.5-3B/blob/main/LICENSE"}, {"title": "Qwen2.5 VL 3B Instruct - API, Providers, Stats | OpenRouter", "url": "https://openrouter.ai/qwen/qwen2.5-vl-3b-instruct:free"}, {"title": "Qwen/Qwen2.5-3B-Instruct · Hugging Face", "url": "https://huggingface.co/Qwen/Qwen2.5-3B-Instruct"}, {"title": "qwen2.5:3b-instruct", "url": "https://ollama.com/library/qwen2.5:3b-instruct"}]

Based on the search results, I found information about the Qwen2.5-3B-Instruct license:

The Qwen2.5-3B-Instruct model uses the qwen-research license, not Apache 2.0. This is confirmed by multiple sources in the Hugging Face repositories.

Additionally, while most Qwen2.5 models are released under the Apache 2.0 license, the 3B and 72B models are under the Qwen license.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. assistant (2026-05-31T04:22:23.155Z)

Both the primary source and independent search results strongly support the claim. Let me verify the actual LICENSE file on the official Qwen HuggingFace repo to independently confirm the license of Qwen2.5-3B-Instruct, since this is the load-bearing fact.

7. user (2026-05-31T04:22:27.522Z)

Based on the web page content, the license tag/metadata listed for this model is:

****qwen-research****

This appears in the page as "License: qwen-research" in the model information section.

8. user (2026-05-31T04:22:28.042Z)

License Analysis

****License Name:**** "Qwen RESEARCH LICENSE AGREEMENT" (verbatim from the document)

****It is NOT Apache 2.0**** ? it is the proprietary Qwen Research License.

****Does it restrict commercial use?**** Yes, explicitly. Section 2.a states:

> "You are granted a non-exclusive, worldwide, non-transferable and royalty-free limited license under Alibaba Cloud's intellectual property or other rights owned by us embodied in the Materials to use, reproduce, distribute, copy, create derivative works of, and make modifications to the Materials ****FOR NON-COMMERCIAL PURPOSES ONLY****."

Section 2.b further clarifies:

> "If you are commercially using the Materials, you shall request a license from us."

This is a restrictive research-only license that prohibits commercial use without explicit permission from Alibaba Cloud.

9. assistant (2026-05-31T04:22:39.361Z)

The evidence is conclusive across multiple independent sources. Let me record my findings.

The claim is verified, not refuted. Walking through the adversarial checklist:

1. ****Supported by quote?**** Yes. The issue author (NaokiAsatani) states verbatim: "the license for Qwen2.5-3B-Instruct is not Apache 2.0, but rather the Qwen RESEARCH LICENSE." The quote directly matches the claim.
2. ****Contradicting evidence?**** None found. Independent WebSearch corroborates: "the 3B and 72B models are under the Qwen license" (not Apache-2.0), while other Qwen2.5 sizes are Apache-2.0.
3. ****Source quality?**** Strong. Beyond the forum post, I independently verified the load-bearing fact at the authoritative source ? the official Qwen HuggingFace repo. The LICENSE file is titled "Qwen RESEARCH LICENSE AGREEMENT" and restricts use "FOR NON-COMMERCIAL PURPOSES ONLY." The model card license metadata tag reads `qwen-research`. This is primary-source confirmation, stronger than the GitHub issue alone.
4. ****Outdated?**** No. The official repo LICENSE remains qwen-research as of this check (2026). The Qwen2.5-3B model has carried the research license since release.
5. ****Marketing/speculation?**** No ? it is a documented license discrepancy backed by the official weights repository.

One nuance for accuracy: the maintainer did not explicitly confirm/refute in the issue (it appears open). But the claim's substance ? that Qwen2.5-3B-Instruct is under the Qwen Research License, not Apache-2.0 ? is independently and authoritatively confirmed regardless of maintainer response. The "disputed" framing is accurate: the issue raises the dispute, and the underlying license fact is verifiably correct. This is directly relevant to muneem's hard constraint #3 (commercially-safe weights), making any InternVL2.5-4B Apache-2.0 attribution unreliable due to its restricted Qwen component.

10. user (2026-05-31T04:22:48.481Z)

Structured output provided successfully

11. assistant (2026-05-31T04:22:52.356Z)

I've completed the verification and submitted my structured output. The claim is ****not refuted**** (confirmed) with high confidence ? the GitHub issue quote matches the claim exactly, and I independently verified the load-bearing fact at the authoritative source (the official Qwen/Qwen2.5-3B-Instruct LICENSE file titled "Qwen RESEARCH LICENSE AGREEMENT" with a non-commercial restriction, plus the `qwen-research` model-card tag).