

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_076beba8-59c\agent-a8e908336cc41a8a8.jsonl`  
- SHA-256: `23eb616fb042d8eb3d028ccbb325254f8b476a981297d1abf296b54802b0ec11`  
- Source modified: `2026-05-31T04:23:37+00:00`  
- Imported at: `2026-07-05T16:48:20+00:00`  
- project: `wf_076beba8-59c`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T04:22:50.749Z)

Adversarial Claim Verifier (voter 1/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Identify the best LOCAL (self-hosted, single NVIDIA CUDA GPU) vision / document-AI / OCR models ? and, as a clearly-flagged approval-gated fallback only, paid cloud API services ? for parsing BANK STATEMENT PDFs into structured transaction tables (date, narration/description, withdrawal/debit, deposit/credit, running balance), for a LOCAL-FIRST personal-finance tool ("muneem").

Target documents: Indian bank statements (HDFC, Axis, ICICI, AU Small Finance) ? dense, multi-column, often SEMI-RULED or borderless tables, sometimes multiple accounts per statement; both born-digital (text layer present) and the harder cases where pdfplumber's extract_tables/coordinate clustering is unreliable (e.g. Axis, whose column layout varies between statements).

HARD CONSTRAINTS (a model that violates these is unusable for us):

1. MUST run 100% LOCALLY for statement contents (highly sensitive). NO cloud APIs by default. Paid/cloud is acceptable ONLY as an explicitly-approved fallback ? still report the best cloud options but label them clearly as "cloud/sensitive ? approval-gated, not default".
2. MUST be usable from Python (transformers / vLLM / Ollama / ONNX Runtime / native lib).
3. License MUST be permissive & commercially safe for BOTH the CODE and the MODEL WEIGHTS: MIT, BSD, Apache-2.0, ISC, HPND. AVOID and explicitly FLAG anything GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, or custom "community"/acceptable-use licenses with restrictions. This is critical ? many document-AI models and their training datasets are non-commercial (e.g. LayoutLMv3 weights are CC-BY-NC; Surya and Marker use GPL-or-commercial dual licensing; several VLM weights ship custom restrictive licenses). Verify the WEIGHTS license, not just the repo license.
4. Hardware: a single consumer/prosumer NVIDIA GPU. Give VRAM tiers (~8-12 GB consumer, ~16-24 GB) and which models/quantizations fit each.

Evaluate and compare these LOCAL approaches ? for each give: code license, MODEL-WEIGHTS license (state explicitly), latest release / maintenance status as of 2025-2026, VRAM need + quantization options, Python integration path, and reliability specifically on DENSE FINANCIAL TABLES / multi-column statements:

- Open vision-language models (VLMs) for document/table understanding: Qwen2-VL and Qwen2.5-VL (note per-size license ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2 / 2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision (flag the Llama Community License restrictions), Moondream2, GOT-OCR2.0, olmOCR and other OCR-specialized VLMs.
- OCR-specialist / document-AI pipelines: PaddleOCR + PP-Structure (table recognition), docTR, EasyOCR, Tesseract (already used), Microsoft Table Transformer / TATR (check code AND PubTables-1M dataset/weights license), Donut, Surya (FLAG license), Marker (FLAG license), MinerU / PDF-Extract-Kit, RapidOCR.

- Table-structure recognition specifically (cell/grid reconstruction for borderless tables): TATR, PP-Structure, UniTable, etc.

Also assess: for SEMI-RULED / borderless bank tables, is a prompt-driven VLM (page image -> JSON transactions) more reliable than a classic OCR + table-structure pipeline? What are the accuracy/hallucination risks of VLMs on financial NUMBERS (transposed digits, dropped/duplicated rows), and how should output be VERIFIED? (muneem already machine-checks parses by balance reconciliation ? opening + sum(deposits) - sum(withdrawals) == closing, plus per-row running-balance walks and cross-checks against the bank's printed summary ? so any model's output can be automatically validated, which should inform the recommendation.)

Finally, briefly cover the best PAID/CLOUD fallbacks (clearly flagged approval-gated for sensitive data): AWS Textract (AnalyzeExpense / Tables), Google Document AI, Azure AI Document Intelligence, Reducto, LlamaParse, Anthropic/OpenAI vision ? note any self-hosting option and data-handling terms.

Deliver a RANKED shortlist of the most promising LOCAL options for muneem (NVIDIA GPU, Python, commercial-license-safe), each with license (code + weights), VRAM tier, maturity, strengths/weaknesses for dense bank-statement tables, and a concrete recommendation: which local model(s) to prototype for muneem's statement parsing and how they'd combine with the existing deterministic parsers + balance-reconciliation verification (e.g. VLM as a fallback/cross-check when deterministic parsing fails to reconcile). Cite sources and call out any license uncertainty explicitly.

Claim under review

"Surya's CODE (the repository LICENSE file at master) is now Apache License 2.0, a permissive license with no revenue restrictions or commercial-use caps. This means the source title's framing of 'GPL-3.0 code' is stale/outdated for the current master branch; only the weights carry the restrictive license."

****Source:**** <https://github.com/VikParuchuri/surya/blob/master/LICENSE> (primary)

****Supporting quote:**** "Apache License
Version 2.0, January 2004
<http://www.apache.org/licenses/>"

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-05-31T04:22:53.255Z)

I'll verify this claim about Surya's code license. Let me check the primary source and look for contradicting evidence.

3. assistant (2026-05-31T04:22:56.274Z)

I'll fetch the primary source (the LICENSE file at master) and search for contradicting evidence in parallel.

4. user (2026-05-31T04:23:01.490Z)

License Analysis

****License Type:**** Apache License 2.0

The file explicitly states at the top: ``"Apache License Version 2.0, January 2004"```

****Additional Terms & Restrictions:****

The license contains standard Apache 2.0 provisions:

- ****Patent Grant:**** Contributors grant perpetual, worldwide, non-exclusive patent licenses, with termination if patent litigation is filed against them
- ****Warranty Disclaimer:**** Work is provided "AS IS" without warranties of any kind
- ****Liability Limitation:**** Contributors are not liable for damages arising from use or inability to use the work
- ****Trademark:**** Does not grant permission to use licensor's trademarks except for reasonable descriptive use

****Commercial Use:**** There are ****no revenue caps or commercial-use limitations****. The license explicitly permits commercial distribution and derivative works, provided conditions (attribution, notice of modifications, license inclusion) are met.

This is a permissive open-source license favorable to commercial applications.

5. user (2026-05-31T04:23:03.018Z)

License Identification

The license stated at the top is: *****"Apache License Version 2.0, January 2004"*****

Additional Clauses

The document contains standard terms but no revenue restrictions or commercial-use limitations. Key provisions include:

- ****Copyright and Patent Grants**:** Contributors grant perpetual, worldwide, royalty-free rights to use, modify, and distribute the work.
- ****Redistribution Requirements**:** Users must provide license copies, mark modifications, and retain attribution notices.
- ****Warranty Disclaimer**:** The work is provided "AS IS" without warranties of any kind.
- ****Liability Limits**:** Contributors are not liable for damages arising from use or inability to use the work.
- ****Trademark Restrictions**:** The license explicitly states it does "not grant permission to use the trade names, trademarks, service marks, or product names of the Licensor."

There are no commercial-use prohibitions?commercial use is permitted, subject to compliance with the stated conditions.

6. user (2026-05-31T04:23:05.288Z)

Web search results for query: "Surya OCR VikParuchuri license GPL Apache 2.0 commercial use"

Links: [{"title":"GitHub - datalab-to/surya: OCR, layout analysis, reading order, table recognition in 90+ languages · GitHub","url":"https://github.com/datalab-to/surya"}, {"title":"surya/LICENSE at master · datalab-to/surya","url":"https://github.com/VikParuchuri/surya/blob/master/LICENSE"}, {"title":"GitHub - CibirQuadance/surya-OCR: OCR and line detection in 90+ languages · GitHub","url":"https://github.com/CibirQuadance/surya-OCR"}, {"title":"Surya OCR | tasarim.ai","url":"https://tasarim.ai/en/models/surya-ocr"}, {"title":"surya-ocr ·"}]

```
PyPI", "url": "https://pypi.org/project/surya-ocr/"], [{"title": "Surya ? Try & Compare | Roboflow Playground", "url": "https://roboflow.com/model/surya"}, {"title": "pitapo/surya · Hugging Face", "url": "https://huggingface.co/pitapo/surya"}, {"title": "View as Markdown", "url": "https://replicate.com/cudanexus/ocr-surya/llms.txt"}, {"title": "surya-ocr - conda-forge | Anaconda.org", "url": "https://anaconda.org/conda-forge/surya-ocr"}, {"title": "surya-ocr 0.2.0", "url": "https://pypi.org/project/surya-ocr/0.2.0/"}]]
```

Based on the search results, here's what I found about Surya OCR's licensing:

Code License

The Surya code is licensed under Apache 2.0.

Model Weights License

The licensing for the model weights is more complex:

The model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue).

Commercial Use

For commercial usage beyond the startup threshold, for broader commercial licensing of the model weights, visit our pricing page here. Additionally, if you want to remove the GPL license requirements for inference or use the weights commercially over the revenue limit, please contact surya@vikas.sh for dual licensing.

Summary

So while the Surya OCR **code** is Apache 2.0 licensed, the **model weights** have restrictions - they're available for free for research and non-commercial use, with a revenue waiver for small startups under \$5M, but commercial use beyond that requires a separate commercial license.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

7. assistant (2026-05-31T04:23:09.506Z)

Both fetches of the LICENSE file at master confirm it is Apache License 2.0 with no commercial restrictions. The web search corroborates this: code is Apache 2.0, weights carry the restrictive license. Let me check the historical/"GPL-3.0 code" framing to confirm whether the claim that it's "stale/outdated" is accurate, and verify the current state isn't a transient.

8. user (2026-05-31T04:23:14.576Z)

Licensing Information from Surya README

Code License:

"The Surya code is licensed under Apache 2.0."

Model Weights License:

The README states the model weights use "a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue)."

Commercial Licensing:

For broader commercial use beyond the restrictions, the README directs users to "visit our pricing page" for commercial licensing options.

GPL:

The README makes no mention of GPL requirements or their removal.

9. user (2026-05-31T04:23:15.467Z)

Web search results for query: "Surya OCR repository GPL-3.0 license change to Apache history 2024"

Links: [{"title": "Surya OCR | tasarim.ai", "url": "https://tasarim.ai/en/models/surya-ocr"}, {"title": "Surya ? Try & Compare | Roboflow Playground", "url": "https://playground.roboflow.com/models/mindee/surya"}, {"title": "surya/LICENSE at master · datalab-to/surya", "url": "https://github.com/VikParuchuri/surya/blob/master/LICENSE"}, {"title": "GitHub - CibirQuadance/surya-OCR: OCR and line detection in 90+ languages · GitHub", "url": "https://github.com/CibirQuadance/surya-OCR"}, {"title": "GitHub - datalab-to/surya: OCR, layout analysis, reading order, table recognition in 90+ languages · GitHub", "url": "https://github.com/datalab-to/surya"}, {"title": "surya-ocr · PyPI", "url": "https://pypi.org/project/surya-ocr/"}, {"title": "surya-ocr - conda-forge | Anaconda.org", "url": "https://anaconda.org/conda-forge/surya-ocr"}, {"title": "GitHub - adithya-s-k/omniparse: Ingest, parse, and optimize any data format ?? from documents to multimedia ?? for enhanced compatibility with GenAI frameworks", "url": "https://github.com/adithya-s-k/omniparse"}, {"title": "Releases · datalab-to/surya", "url": "https://github.com/datalab-to/surya/releases"}, {"title": "surya-ocr 0.1.2", "url": "https://pypi.org/project/surya-ocr/0.1.2/"}]

Based on the search results, here's what I found regarding Surya OCR's license change:

License Change History

Surya was first released in January 2024 under the GPL-3.0 license. However, there has been a significant license change since then.

The Surya code is now licensed under Apache 2.0, while the model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue).

The current official repository at GitHub (datalab-to/surya) reflects this Apache 2.0 license for the code, representing a shift from the original GPL-3.0 license that was in place at the project's initial release. This change makes the code more permissive for commercial use while maintaining restrictions on the model weights themselves through a separate licensing scheme.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

10. assistant (2026-05-31T04:23:27.634Z)

I have strong, consistent evidence from multiple sources. Let me assess the claim against the checklist.

Findings:

- Primary source confirms the quote.** Two independent fetches of the LICENSE file at `master` (both the GitHub blob view and the raw githubusercontent URL) confirm the file's top reads "Apache License Version 2.0, January 2004" with standard Apache 2.0 terms and explicitly no revenue caps or commercial-use limitations. The supporting quote is accurate and not cherry-picked.
- The "stale GPL-3.0" framing is verified accurate.** The web search confirms the historical trajectory: Surya was first released in January 2024 under GPL-3.0, then the code was relicensed to Apache 2.0. So the claim that "GPL-3.0 code" framing is outdated for the current master branch is correct ? it reflects the project's initial state, not current master.
- The weights-only restriction is confirmed.** The README (datalab-to/surya, the current official repo) states "The Surya code is licensed under Apache 2.0" while the weights use "a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue)." This directly supports the claim's final clause: "only the weights carry the restrictive license."

****One nuance to note (does not refute):**** The original source's "GPL-or-commercial dual licensing" framing wasn't purely wrong ? historically there *was* a GPL dimension (the search result mentions "if you want to remove the GPL license requirements for inference"). But the claim under review is narrowly about the ****CODE**** repository LICENSE file at master being Apache 2.0, and the framing of "GPL-3.0 code" being stale. That narrow claim is well-supported. The weights' actual license is "modified AI Pubs Open RAIL-M" (still restrictive/non-commercial above threshold), which the claim correctly characterizes as "the restrictive license" without misstating it as permissive.

****Checklist verdict:****

- Supported by quote: YES (verified at master, twice).
- Contradicting evidence: NONE found; corroborated by README and search.
- Source quality: PRIMARY source (the actual LICENSE file) ? appropriate for this claim.
- Outdated: NO ? this is the *current* master state; the claim correctly flags the *opposing* "GPL code" framing as the outdated one.
- Marketing/speculation: NO ? direct file inspection.

The claim survives. It is well-supported, current, and sourced primarily.

11. user (2026-05-31T04:23:34.182Z)

Structured output provided successfully

12. assistant (2026-05-31T04:23:37.291Z)

The claim survives verification. I confirmed via the primary source (the actual LICENSE file at master, fetched two ways) that Surya's code is Apache 2.0 with no commercial restrictions, that the "GPL-3.0 code" framing reflects the project's January 2024 initial state and is now stale, and that only the model weights carry the restrictive (modified AI Pubs Open RAIL-M) license. Structured output submitted with refuted=false, high confidence.