

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\ a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_89367d2c-824\agent-a43b077ec459a8f71.jsonl`  
- SHA-256: `8f78db429e0f435d4b80f9cd94606754775e32968b0d575bd09671484433bef3`  
- Source modified: `2026-05-31T08:36:01+00:00`  
- Imported at: `2026-07-05T16:48:21+00:00`  
- project: `wf_89367d2c-824`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T08:33:52.116Z)

Context: "muneem" is a LOCAL-FIRST personal-finance tool. Goal: turn BANK STATEMENT PDF pages (Indian banks ? HDFC, Axis, ICICI, AU Small Finance) into structured transactions (txn_date, narration, debit/withdrawal, credit/deposit, balance). It will be used as an L3 FALLBACK that only fires when deterministic parsing fails a balance-reconciliation check ? so a probabilistic model is acceptable ONLY because its output is machine-verified (opening + ?deposits ? ?withdrawals == closing).

HARD CONSTRAINTS (a model that violates these is unusable):

1) Must run 100% LOCALLY for statement contents (sensitive); cloud only as a flagged, approval-gated fallback.

2) Python-usable (transformers / vLLM / Ollama / ONNX).

3) License must be commercially safe for BOTH the CODE and the MODEL WEIGHTS:

MIT/BSD/Apache-2.0/ISC/HPND OK; FLAG and avoid GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, and custom "community"/acceptable-use licenses. Known traps to verify: LayoutLMv3 weights are CC-BY-NC; Surya & Marker are GPL-or-commercial dual; Llama-3.2-Vision uses the Llama Community License. CHECK THE WEIGHTS LICENSE, not just the repo.

4) Single consumer/prosumer NVIDIA GPU; report VRAM tiers (8/12/16/24 GB).

Research THIS angle with web search, and fetch key pages (model cards, LICENSE files, docs) to confirm license + spec details:

Open vision-language models (VLMs) for document/table understanding (2025-2026). For EACH of: Qwen2-VL and Qwen2.5-VL (state license PER SIZE ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2/InternVL2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision, Moondream2, GOT-OCR2.0, olmOCR ? give the MODEL-WEIGHTS license EXPLICITLY (commercially usable yes/no), parameter count + VRAM, and the local Python run path (transformers/vLLM/Ollama). Flag any non-commercial or custom-restricted weights.

Return a concise findings brief: bullet points stating the specific facts (CODE license, WEIGHTS license ? quote the exact license name, VRAM, latest-release date, reliability) each followed by its SOURCE URL. Where a license or fact is uncertain or you couldn't confirm it, say so explicitly rather than guessing.

2. assistant (2026-05-31T08:33:54.885Z)

I'll search for the latest open VLM models suitable for document/table understanding, focusing on license compliance and technical specs.

3. user (2026-05-31T08:34:01.017Z)

Web search results for query: "Qwen2-VL Qwen2.5-VL model weights license Apache 2024 2025"

Links: [{"title": "Qwen - Wikipedia", "url": "https://en.wikipedia.org/wiki/Qwen"}, {"title": "License for the model weights · Issue #963 · QwenLM/Qwen3-VL", "url": "https://github.com/QwenLM/Qwen3-VL/issues/963"}, {"title": "LICENSE · Qwen/Qwen2.5-72B-Instruct at main", "url": "https://huggingface.co/Qwen/Qwen2.5-72B-"}]

```
Instruct/blob/main/LICENSE"},"{"title":"Qwen2.5: A Party of Foundation Models! | Qwen", "url":"https://qwenlm.github.io/blog/qwen2.5/"},"{"title":"Qwen (Tongyi Qianwen): Alibaba's Open-Weight AI Model Family", "url":"https://skywork.ai/blog/qwen-tongyi-qianwen-open-weight-ai-model/"},"{"title":"Qwen2.5 VL 7B Instruct ? Try & Compare | Roboflow Playground", "url":"https://playground.roboflow.com/models/qwen/qwen2-5-vl-7b-instruct"},"{"title":"LICENSE · Qwen/Qwen2.5-7B at main", "url":"https://huggingface.co/Qwen/Qwen2.5-7B/blob/main/LICENSE"},"{"title":"Qwen Model Series: Open-Source Multimodal LLMs", "url":"https://www.emergentmind.com/topics/qwen-model-series"},"{"title":"Qwen2.5 Model Family Overview", "url":"https://www.emergentmind.com/topics/qwen2-5-model-family"}]
```

Based on the search results, here's what I found regarding the Qwen2-VL and Qwen2.5-VL model weights and licensing:

Qwen2-VL

Alibaba released Qwen2-VL with variants of 2 billion and 7 billion parameters.

Qwen2.5-VL

In January 2025, Qwen2.5-VL was released with variants of 3, 7, 32, and 72 billion parameters.

License Information

The licensing structure for these models varies by size variant:

All models except the 72B variant are licensed under the Apache 2.0 license. More specifically, Qwen2.5-VL-7B-Instruct is released under the Apache 2.0 license.

However, there appears to be some variation in licensing for other variants. All open-source models, except for the 3B and 72B variants, are licensed under Apache 2.0. This indicates that the 3B and 72B variants may use a different license (the Qwen License), while the 7B variant uses the more permissive Apache 2.0 license.

The model weights are publicly available on platforms like Hugging Face and GitHub.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

4. user (2026-05-31T08:34:01.351Z)

Web search results for query: "Microsoft Phi-3.5-vision model weights license commercial 2025"

Links: [{"title":"phi-3.5-vision-instruct Model by Microsoft | NVIDIA NIM", "url":"https://build.nvidia.com/microsoft/phi-3_5-vision-instruct/modelcard"}, {"title":"Microsoft releases powerful new Phi-3.5 models, beating Google, OpenAI and more | VentureBeat", "url":"https://venturebeat.com/ai/microsoft-releases-powerful-new-phi-3-5-models-beating-google-openai-and-more"}, {"title":"Phi-3.5-vision instruct (128k)", "url":"https://ai.azure.com/catalog/models/Phi-3.5-vision-instruct"}, {"title":"Microsoft releases new Phi 3.5 open-source language and vision models", "url":"https://the-decoder.com/microsoft-releases-new-phi-3-5-open-source-language-and-vision-models/"}, {"title":"microsoft / phi-3.5-mini", "url":"https://docs.api.nvidia.com/nim/reference/microsoft-phi-3_5-mini"}, {"title":"microsoft/Phi-3-vision-128k-instruct · Hugging Face", "url":"https://huggingface.co/microsoft/Phi-3-vision-128k-instruct"}, {"title":"Phi-3 Technical Report: A Highly Capable Language Model Locally on Your Phone - Microsoft Research", "url":"https://www.microsoft.com/en-us/research/publication/phi-3-technical-report-a-highly-capable-language-model-locally-on-your-phone/"}, {"title":"Microsoft AI Releases Phi 3.5 mini, MoE and Vision with 128K context, Multilingual and MIT License - MarkTechPost", "url":"https://www.marktechpost.com/2024/08/21/microsoft-ai-releases-phi-3-5-mini-moe-and-vision-with-128k-context-multilingual-and-mit-license/"}, {"title":"microsoft/Phi-3.5-vision-instruct · Hugging Face", "url":"https://huggingface.co/microsoft/Phi-3.5-vision-instruct"}, {"title":"Phi Open Models - Small Language Models | Microsoft Azure", "url":"https://azure.microsoft.com/en-us/products/phi"}]

Based on the search results, here's what I found about the Microsoft Phi-3.5-vision model weights and licensing:

License

The model is licensed under the MIT license. All three models are available for developers to download, use, and fine-tune/customize on Hugging Face under a Microsoft-branded MIT License that allows for commercial usage and modification without restrictions.

Commercial Use

The model is intended for broad commercial and research use in English. This license allows developers to freely use, modify, merge, publish, distribute, sublicense, or sell copies of the software. The license also includes a disclaimer that the software is provided "as is," without warranties of any kind. Microsoft and other copyright holders are not liable for any claims, damages, or other liabilities that may arise from the software's use.

Availability

The Phi 3.5 models are available under the MIT license on Hugging Face and through Microsoft's Azure AI Studio.

In summary, the Microsoft Phi-3.5-vision model weights are released under a permissive MIT license that explicitly allows for commercial use without restrictions.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

5. user (2026-05-31T08:34:01.441Z)

Web search results for query: "InternVL2 InternVL2.5 model weights license Apache MIT 2025"

Links: [{"title": "OpenGVLab/InternVL2_5-Pretrain-Models · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL2_5-Pretrain-Models"}, {"title": "OpenGVLab/InternVL3_5-8B · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL3_5-8B"}, {"title": "About the license of internVL2.5-4B model · Issue #900 · OpenGVLab/InternVL", "url": "https://github.com/opengvlab/internvl/issues/900"}, {"title": "Quick Start of InternVL 2.5 Series ? InternVL", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/quick_start.html"}, {"title": "Expanding Performance Boundaries of Open-Source Multimodal Models with Model, Data, and Test-Time Scaling", "url": "https://arxiv.org/pdf/2412.05271"}, {"title": "InternVL2.5", "url": "https://internvl.github.io/blog/2024-12-05-InternVL-2.5/"}, {"title": "GitHub - OpenGVLab/InternVL: [CVPR 2024 Oral] InternVL Family: A Pioneering Open-Source Alternative to GPT-4o. ?GPT-4o????????????", "url": "https://github.com/opengvlab/internvl"}, {"title": "OpenGVLab/InternVL2-Pretrain-Models · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL2-Pretrain-Models"}, {"title": "OpenGVLab/InternVL2_5-2B · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL2_5-2B"}]

Based on the search results, here's what I found regarding InternVL2 and InternVL2.5 model weights licensing:

InternVL2.5 License

InternVL 2.5 is released under the MIT license. However, it's important to note that InternVL2.5 uses pre-trained components (such as Qwen) which are licensed under the Apache License 2.0.

InternVL2 License

InternVL 2.0 is released under the MIT license, though it also uses pre-trained components that may have different licenses.

Important Licensing Consideration

There was a notable licensing discrepancy reported: The license for some Qwen components is not Apache 2.0 but the Qwen RESEARCH LICENSE, which prohibits commercial use. This is an important consideration if you're planning commercial deployment of these models.

The model weights themselves are released under permissive licenses (MIT for the InternVL project), but users should verify the licenses of all component models used in their specific configuration before deploying commercially.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. user (2026-05-31T08:34:01.886Z)

Web search results for query: "MiniCPM-V 2.6 model weights license commercially usable 2025"

Links: [{"title": "openbmb/MiniCPM-V-2_6 · Hugging Face", "url": "https://huggingface.co/openbmb/MiniCPM-V-2_6"}, {"title": "openbmb/MiniCPM-V · Hugging Face", "url": "https://huggingface.co/openbmb/MiniCPM-V"}, {"title": "minicpm-v:8b-2.6-q4_0/license", "url": "https://ollama.com/library/minicpm-v:8b-2.6-q4_0/blobs/8603ca877636"}, {"title": "GitHub - QDottech/MiniCPM-V: MiniCPM-V 2.6: A GPT-4V Level MLLM for Single Image, Multi Image and Video on Your Phone · GitHub", "url": "https://github.com/QDottech/MiniCPM-V"}, {"title": "openbmb/MiniCPM-o-2_6 · Hugging Face", "url": "https://huggingface.co/openbmb/MiniCPM-o-2_6"}, {"title": "openbmb/MiniCPM-V-4.6 · Hugging Face", "url": "https://huggingface.co/openbmb/MiniCPM-V-4.6"}, {"title": "minicpm-v:8b-2.6-q3_K_S/license", "url": "https://www.ollama.com/library/minicpm-v:8b-2.6-q3_K_S/blobs/8603ca877636"}, {"title": "minicpm-v/latest/blobs/8603ca877636"}, {"title": "GitHub - OpenBMB/MiniCPM-V: A Pocket-Sized MLLM for Ultra-Efficient Image and Video Understanding on Your Phone · GitHub", "url": "https://github.com/OpenBMB/MiniCPM-V"}, {"title": "openbmb/MiniCPM-V-2 · Hugging Face", "url": "https://huggingface.co/openbmb/MiniCPM-V-2"}]

Based on the search results, here's what I found regarding the MiniCPM-V 2.6 model weights and commercial usability:

License Information

The models and weights of MiniCPM are completely free for academic research. After filling out a "questionnaire" for registration, MiniCPM-V 2.6 weights are also available for free commercial use.

License Terms

Community use of the MiniCPM series models must comply with Apache 2.0 and the "MiniCPM Model Community License Agreement."

Commercial Usage Conditions

For commercial use, there are specific conditions:

If you or your affiliates' services or products deploy the model on edge-side devices not exceeding 5,000 units, or provide applications with a daily active user count (DAU) of less than 1 million, you can apply to OpenBMB for permission and, after filling out the registration questionnaire, may be allowed to use it commercially for free.

Otherwise, please email (cpm@modelbest.cn) to apply for authorization from OpenBMB, which may, at its discretion, grant permission.

In summary, MiniCPM-V 2.6 model weights are commercially usable, but there are thresholds and

requirements involved?requiring registration and questionnaire completion, with larger-scale deployments requiring explicit authorization from OpenBMB.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

7. assistant (2026-05-31T08:34:03.272Z)

Now let me search for the remaining models:

8. user (2026-05-31T08:34:09.236Z)

Web search results for query: "Moondream2 model weights license Apache MIT 2024 2025"

Links: [{"title":"Moondream 2: Multimodal and Vision Analysis","url":"https://blog.roboflow.com/moondream-2/"},"{"title":"Overview | Moondream Docs","url":"https://docs.moondream.ai/"},"{"title":"Moondream 2 ? Try & Compare | Roboflow Playground","url":"https://roboflow.com/model/moondream-2"},"{"title":"vikhyatk/moondream2 · Hugging Face","url":"https://huggingface.co/vikhyatk/moondream2"},"{"title":"README.md · vikhyatk/moondream2 at 2025-06-21","url":"https://huggingface.co/vikhyatk/moondream2/blob/2025-06-21/README.md"},"{"title":"Finetuning Moondream2 for Computer Vision Tasks","url":"https://blog.roboflow.com/finetuning-moondream2/"},"{"title":"Introducing Moondream 0.5B: The World's Smallest Vision-Language Model. | Moondream","url":"https://moondream.ai/blog/introducing-moondream-0-5b"},"{"title":"Moondream2 by vikhyatk ? VRAM 3.9GB | LLM Explorer","url":"https://llm-explorer.com/model/vikhyatk%2Fmoondream2,4ZvnS9KBBi7BV7u3f5Snag"},"{"title":"moondream:v2","url":"https://ollama.com/library/moondream:v2"},"{"title":"vikhyatk/moondream2 at main","url":"https://huggingface.co/vikhyatk/moondream2/tree/main"}]

Based on the search results, here's what I found about the Moondream2 model weights license:

License Information

The model is licensed under an Apache 2.0 license. This is confirmed across multiple sources, including the official Hugging Face repository where the license is listed as apache-2.0.

Moondream2 is open-source under the Apache 2.0 license, allowing commercial use as well.

Model Timeline

The model was first released in early 2024 and updated through mid-2025. The most recent release appears to be from 2025-06-21.

Important Note

It's worth noting that there appears to be a distinction between different distributions of Moondream2. While the open-source version on Hugging Face (by vikhyatk) uses the Apache 2.0 license, newer model weights are licensed under the Business Source License 1.1 with an Additional Use Grant (No Third-Party Service) according to the official Moondream documentation, suggesting licensing may vary depending on which distribution or version you use.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

9. user (2026-05-31T08:34:09.691Z)

Web search results for query: "Mistral Pixtral model weights license commercial 2025"

Links: [{"title":"[Deprecated] Pixtral Large | Mistral AI","url":"https://mistral.ai/news/pixtral-large/"},"{"title":"Pixtral Large: Mistral AI's Groundbreaking Leap in Large Language Models | Encord","url":"https://encord.com/blog/pixtral-

```
large-explained/"},"title":"mistralai/Pixtral-12B-2409 · Hugging
Face","url":"https://huggingface.co/mistralai/Pixtral-12B-2409"},"title":"Mistral closes in on
Big AI rivals with new open-weight frontier and small models |
TechCrunch","url":"https://techcrunch.com/2025/12/02/mistral-closes-in-on-big-ai-rivals-with-
mistral-3-open-weight-frontier-and-small-models/"},"title":"What is Mistral AI? |
IBM","url":"https://www.ibm.com/think/topics/mistral-ai"},"title":"Model weights | Mistral
AI","url":"https://platform-docs-public.pages.dev/getting-
started/models/weights/"},"title":"Mistral Models
Overview","url":"https://docs.mistral.ai/models/overview"},"title":"Models - from cloud to edge
| Mistral","url":"https://mistral.ai/models/"},"title":"What is Mistral AI? Features, Pricing,
and Use Cases","url":"https://www.walturn.com/insights/what-is-mistral-ai-features-pricing-and-
use-cases"}]]
```

Based on the search results, here's what I found about Mistral Pixtral model weights and licensing for commercial use in 2025:

Pixtral Model Variants and Licenses

****Pixtral 12B:****

Pixtral 12B is an open multimodal model, offered under an Apache 2.0 license, which permits their use without any constraints.

****Pixtral Large:****

Pixtral Large is under Mistral Research License. More specifically, the model is available under the Mistral Research License (MRL) for research and educational use; and the Mistral Commercial License for experimentation, testing, and production for commercial purposes.

Commercial Use

For Pixtral Large's commercial deployment, commercial licenses can be granted on request through contacting Mistral directly.

Summary

- ****Pixtral 12B****: Fully open under Apache 2.0 with no commercial restrictions
- ****Pixtral Large****: Requires either the Mistral Research License (non-commercial) or a separate commercial license from Mistral for production use

The key distinction is that the smaller 12B model is freely available for any use, while the larger 124B model requires explicit commercial licensing from Mistral for production deployments.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

10. user (2026-05-31T08:34:12.435Z)

Web search results for query: "Meta Llama-3.2-Vision model weights license community 2025"

```
Links: [{"title":"meta-llama/Llama-3.2-11B-Vision · Hugging
Face","url":"https://huggingface.co/meta-llama/Llama-3.2-11B-Vision"},"title":"Meta AI ? All
Models Available for Consumers and Developers in
2025","url":"https://www.datastudios.org/post/meta-ai-all-models-available-for-consumers-and-
developers-in-2025"},"title":"Llama 3.2 Overview ? NVIDIA NIM for Vision Language Models
(VLMs)","url":"https://docs.nvidia.com/nim/vision-language-
models/1.2.0/examples/llama3-2/overview.html"},"title":"llama-models/models/llama3_2/LICENSE at
main · meta-llama/llama-models","url":"https://github.com/meta-llama/llama-
models/blob/main/models/llama3_2/LICENSE"},"title":"Llama 3.2 Complete Guide: Lightweight &
Vision Models","url":"https://llamaimodel.com/download-3-2/"},"title":"Llama 3.2:
Revolutionizing edge AI and vision with open, customizable
models","url":"https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-
devices/"},"title":"meta-llama/Llama-3.2-90B-Vision · Hugging
```

```
Face", "url": "https://huggingface.co/meta-llama/Llama-3.2-90B-Vision"}, {"title": "Meta Llama 3 License", "url": "https://www.llama.com/llama3/license/"}, {"title": "llama3.2-vision/license", "url": "https://ollama.com/library/llama3.2-vision:latest/blobs/d3ed60b91757"}, {"title": "GitHub - meta-llama/llama-models: Utilities intended for use with Llama models. · GitHub", "url": "https://github.com/meta-llama/llama-models"}]}
```

Based on the search results, here's what I found about Meta's Llama-3.2-Vision model weights, license, and community availability:

Model Weights and Availability

The Meta Llama 3.2 Vision collection of multimodal large language models is available in 11B and 90B sizes, and models are available for download on Hugging Face and partner platforms including AMD, AWS, Databricks, Dell, Google Cloud, Groq, IBM, Intel, Microsoft Azure, NVIDIA, Oracle Cloud, and Snowflake.

License Terms

Use of Llama 3.2 is governed by the Llama 3.2 Community License, a custom, commercial license agreement. Llama models are released under Meta's community licenses as open-weight with acceptable-use restrictions, and while they're not OSI-open-source, they are broadly usable for commercial apps within license terms.

The Community License is free for research and most businesses; companies with over 700 million monthly active users need a Meta-approved commercial deal.

Key Restrictions and Requirements

If you distribute the Llama Materials or create derivative works, you must provide a copy of the license agreement and prominently display "Built with Llama," and if you use the outputs to create or improve an AI model, you must include "Llama" at the beginning of the model name.

Vision checkpoints are not legally distributed in the European Union.

Community Features

The Llama 3.2 model collection supports the ability to leverage outputs to improve other models including synthetic data generation and distillation, and the Llama 3.2 Community License allows for these use cases.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

11. user (2026-05-31T08:34:14.854Z)

Web search results for query: "GOT-OCR2.0 model weights license commercial 2025"

Links: [{"title": "GOT-OCR2 · Hugging Face", "url": "https://huggingface.co/docs/transformers/model_doc/got_ocr2"}, {"title": "GitHub - Ucas-HaoranWei/GOT-OCR2.0: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model · GitHub", "url": "https://github.com/Ucas-HaoranWei/GOT-OCR2.0/"}, {"title": "abhinand/GOT-OCR-2.0-unofficial · Hugging Face", "url": "https://huggingface.co/abhinand/GOT-OCR-2.0-unofficial"}, {"title": "GOT-OCR2.0 | Ecosystem Directory | market.dev", "url": "https://explore.market.dev/ecosystems/python/projects/got-ocr2-0"}, {"title": "GOT-OCR2_0", "url": "https://www.promptlayer.com/models/got-ocr20"}, {"title": "GOT-OCR-2.0 huggingface.co api & abhinand GOT-OCR-2.0 github AI Model - Toolify", "url": "https://www.toolify.ai/ai-model/abhinand-got-ocr-2-0"}, {"title": "GOT-OCR2_0 | AI Model Details", "url": "https://www.aimodels.fyi/models/huggingFace/got-ocr20-ucaslcl"}, {"title": "GitHub - shaneholloman/got-ocr2.0: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end"}]

```
Model", "url": "https://github.com/shaneholloman/got-ocr2.0"}, {"title": "GitHub - MacOS/got-ocr2.0: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model", "url": "https://github.com/MacOS/got-ocr2.0"}, {"title": "GOT OCR 2.0 (Vision) API on fal", "url": "https://fal.ai/models/fal-ai/got-ocr/v2"}]
```

Based on the search results, here's what I found regarding GOT-OCR2.0 model weights and licensing:

License Information

The data, code, and checkpoint are intended and licensed for research use only. They are also restricted to use that follow the license agreement of Vary.

This means that **the GOT-OCR2.0 model weights are NOT licensed for commercial use** according to the official documentation. The license is explicitly restricted to research purposes only, and users must also comply with the Vary model's license agreement (which GOT-OCR2.0 is built upon).

Model Release Timeline

The model was released on 2024-09-03 and added to Hugging Face Transformers on 2025-01-31. The model was open-sourced with its codes and weights on September 3rd, 2024.

Key Takeaway for Commercial Use

If you're interested in using GOT-OCR2.0 for commercial purposes, you would need to contact the creators or seek an alternative license arrangement, as the current open-source distribution is restricted to research use only.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

12. assistant (2026-05-31T08:34:16.940Z)

Now let me search for olmOCR and get detailed specs (VRAM, params, release dates):

13. user (2026-05-31T08:34:22.472Z)

Web search results for query: "olmOCR model weights license commercial 2025"

Links: [{"title": "allenai/olmOCR-2-7B-1025 · Hugging Face", "url": "https://huggingface.co/allenai/olmOCR-2-7B-1025"}, {"title": "GitHub - datalab-to/surya: OCR, layout analysis, reading order, table recognition in 90+ languages · GitHub", "url": "https://github.com/datalab-to/surya"}, {"title": "olmOCR-2-7B-1025 OCR Model", "url": "https://www.emergentmind.com/topics/olmocr-2-7b-1025"}, {"title": "olmOCR 2: Unit test rewards for document OCR | Ai2", "url": "https://allenai.org/blog/olmocr-2"}, {"title": "Meet olmOCR - Open Source OCR for Accurate Document Conversion - Abdul Aziz Ahwan", "url": "https://www.abdulazizahwan.com/2025/03/meet-olmocr-open-source-ocr-for-accurate-document-conversion.html"}, {"title": "olmOCR 2", "url": "https://lmstudio.ai/models/olmocr2"}, {"title": "olmOCR: Unlocking Trillions of Tokens in PDFs with Vision Language Models", "url": "https://olmocr.allenai.org/papers/olmocr_3pg_sample.pdf"}, {"title": "olmOCR Open Source OCR System for AI Training Using PDFs & Documents", "url": "https://www.geeky-gadgets.com/olmocr-ocr-model-for-ai-training/"}, {"title": "GitHub - allenai/olmocr: Toolkit for linearizing PDFs for LLM datasets/training · GitHub", "url": "https://github.com/allenai/olmocr"}, {"title": "(PDF) olmOCR 2: Unit Test Rewards for Document OCR", "url": "https://www.researchgate.net/publication/396790587_olmOCR_2_Unit_Test_Rewards_for_Document_OCR"}]

Based on the search results, here's what I found about olmOCR model weights licensing:

License Information

The olmOCR model weights are licensed under Apache 2.0. It is intended for research and educational use in accordance with AI2's Responsible Use Guidelines.

Commercial Use

The complete suite—including model weights, training/inference code, and both the supervised (olmOCR-mix-1025) and synthetic (olmOCR2-synthmix-1025) datasets—is released under permissive open-source licenses (e.g., Apache 2.0, MIT), with explicit notation on API or usage restrictions as appropriate.

The Apache 2.0 license is a permissive open-source license that generally allows commercial use, modification, and distribution, though it requires attribution and includes liability disclaimers. However, the explicit statement that the model is "intended for research and educational use in accordance with AI2's Responsible Use Guidelines" suggests there may be additional restrictions beyond the formal license terms that users should be aware of.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

14. user (2026-05-31T08:34:23.553Z)

Web search results for query: "InternVL2.5 parameters VRAM memory requirements 2B 4B 8B"

Links: [{"title": "Deep Learning VRAM Calculator - GPU Memory Requirements for LLMs & AI | DataZier", "url": "https://www.datazier.com/free-tools/vram-calculator/"}, {"title": "OpenGVLab/Mini-InternVL-Chat-4B-V1-5 · How much VRam does it need?", "url": "https://huggingface.co/OpenGVLab/Mini-InternVL-Chat-4B-V1-5/discussions/6"}, {"title": "Will This LLM Fit My GPU? VRAM Requirements for Every Model Size ? ai.rs", "url": "https://ai.rs/ai-developer/will-llm-fit-my-gpu-vram-requirements"}, {"title": "GitHub - OpenGVLab/InternVL: [CVPR 2024 Oral] InternVL Family: A Pioneering Open-Source Alternative to GPT-4o. ??GPT-4o???????????? .", "url": "https://github.com/OpenGVLab/InternVL"}, {"title": "LLM VRAM Calculator - Advanced GPU Memory Calculator for Large Language Models", "url": "https://llmvrascalculator.com/"}, {"title": "InternVL2.5", "url": "https://internvl.github.io/blog/2024-12-05-InternVL-2.5/"}, {"title": "Expanding Performance Boundaries of Open-Source Multimodal Models with Model, Data, and Test-Time Scaling", "url": "https://arxiv.org/pdf/2412.05271"}, {"title": "VRAM Calculator - Check GPU Memory Requirements for Gaming & AI | Orbit2x", "url": "https://orbit2x.com/vram-calculator"}]

Based on the search results, I found a reference to InternVL2 documentation that contains memory usage information. According to a discussion on Hugging Face, detailed memory usage testing information is available in the InternVL2.0 deployment documentation, which should cover the InternVL2.5 series as well since they share the same core architecture.

However, the search results don't directly display the specific VRAM requirements for the 2B, 4B, and 8B variants of InternVL2.5. The documentation referenced (at <https://internvl.readthedocs.io/en/latest/internvl2.0/deployment.html#memory-usage-testing>) would contain this information, but the actual memory specifications aren't included in the search results returned.

InternVL2.5 is an advanced multimodal large language model series with parameter coverage ranging from 1B to 78B, confirming these model sizes exist, but the specific VRAM requirements for each size would need to be obtained directly from the official InternVL documentation linked above.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

15. user (2026-05-31T08:34:24.264Z)

Web search results for query: "MiniCPM-V 2.6 parameters 8B VRAM memory specifications"

Links: [{"title": "MiniCPM-V-2.6 ? Xinference", "url": "https://inference.readthedocs.io/en/stable/models/builtin/llm/minicpm-v-2.6.html"}, {"title": "minicpm-v:8b", "url": "https://ollama.com/library/minicpm-v:8b"}, {"title": "minicpm-v:8b-2.6-fp16", "url": "https://ollama.com/library/minicpm-v:8b-2.6-fp16"}, {"title": "GitHub - QAdottech/MiniCPM-V: MiniCPM-V 2.6: A GPT-4V Level MLLM for Single Image, Multi Image and Video on Your Phone · GitHub", "url": "https://github.com/QAdottech/MiniCPM-V"}, {"title": "MiniCPM-V 2.6 | Radxa Docs", "url": "https://docs.radxa.com/en/orion/o6/app-development/artificial-intelligence/Multimodality/minicpm-v-2-6"}, {"title": "minicpm-v:8b-2.6-q8_0", "url": "https://ollama.com/library/minicpm-v:8b-2.6-q8_0"}, {"title": "openbmb/MiniCPM-V-2_6 · Hugging Face", "url": "https://huggingface.co/openbmb/MiniCPM-V-2_6"}, {"title": "openbmb/minicpm-v2.6", "url": "https://ollama.com/openbmb/minicpm-v2.6"}, {"title": "minicpm-v:8b-2.6-q4_K_M", "url": "https://ollama.com/library/minicpm-v:8b-2.6-q4_K_M"}, {"title": "MiniCPM-V Series | OpenBMB/MiniCPM-o | DeepWiki", "url": "https://deepwiki.com/OpenBMB/MiniCPM-o/1.2-minicpm-v-series"}]

Based on the web search results, here are the key specifications for MiniCPM-V 2.6:

Parameters and Architecture

MiniCPM-V 2.6 is built on SigLip-400M and Qwen2-7B with a total of 8B parameters.

VRAM and Memory Requirements

Regarding memory specifications:

- The int4 quantized version requires lower GPU memory (7GB) usage
- You can run the model on multiple low VRAM GPUs (12 GB or 16 GB) by distributing the model's layers across multiple GPUs

Key Efficiency Features

MiniCPM-V 2.6 shows state-of-the-art token density (number of pixels encoded into each visual token), producing only 640 tokens when processing a 1.8M pixel image, which is 75% fewer than most models. This directly improves the inference speed, first-token latency, memory usage, and power consumption.

The model supports various deployment options including llama.cpp and ollama support for efficient CPU inference on local devices, int4 and GGUF format quantized models in 16 sizes, and vLLM support for high-throughput and memory-efficient inference.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

16. user (2026-05-31T08:34:25.685Z)

Web search results for query: "Qwen2.5-VL parameters VRAM requirements 7B 32B specifications"

Links: [{"title": "Qwen/Qwen2.5-VL-7B-Instruct · Hardware and vram requiremnt to run this model?", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct/discussions/18"}, {"title": "Qwen/Qwen2.5-VL-72B-Instruct · Model Memory Requirements", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-72B-Instruct/discussions/3"}, {"title": "System Requirements | QwenLM/Qwen2.5-VL | DeepWiki", "url": "https://deepwiki.com/QwenLM/Qwen2.5-VL/1.2-system-requirements"}, {"title": "Qwen2.5 VL 72B VRAM Needs: Why Vision Models ... - Novita", "url": "https://blogs.novita.ai/qwen2-5-vl-72b-vram/"}, {"title": "Qwen VRAM Requirements (All Model Sizes) GIGAGPU", "url": "https://gigagpu.com/qwen-vram-requirements/"}, {"title": "Qwen2.5-Coder 14B VRAM Requirements ? Q4, Q5, Q8, FP16 Hardware Guide | Will It Run AI Blog", "url": "https://willitrunai.com/blog/qwen-2-5-coder-14b-vram-requirements/"}, {"title": "Qwen 2.5 7B VRAM Tips Every Dev Should Know | by Novita AI | Medium", "url": "https://medium.com/@marketing_novita.ai/qwen-2-5-7b-vram-tips-every-dev-should-know-932303373ff0"}, {"title": "Qwen 2.5 VRAM Requirements: 9B, 27B, 35B, 72B [2026 Guide] | CraftRigs", "url": "https://craftrigs.com/guides/qwen-35-local-llm-hardware-"}]

```
requirements/"},"title":"Promptquorum","url":"https://www.promptquorum.com/local-llms/qwen-local-deployment-guide-2026"]]
```

Based on the search results, here's what I found about Qwen2.5-VL specifications and VRAM requirements:

Model Sizes and Parameters

Qwen2.5-VL is available in multiple sizes (3B, 7B, 32B, and 72B parameters), each with different resource requirements.

VRAM Requirements

The VRAM requirements vary significantly by model size and precision:

****For the 72B model:****

72 billion parameters require 144 GB of VRAM if running in bf16, 72 GB for fp8, and 36 GB for 4-bit datatype. Deploying Qwen2.5-VL-72B locally can require up to 384 GB of VRAM.

****For the 7B model:****

The Qwen2.5-VL-7B model runs without any issues with 24GB VRAM (like an RTX 4090) when using flash_attention_2.

Important Considerations

Actual memory usage is typically at least 1.2 times higher than theoretical minimums. Additionally, additional memory may be required when processing high-resolution images, multiple images in a single request, video inputs, and batch processing.

Vision-language models (Qwen2-VL) require slightly more VRAM than text-only equivalents due to the vision encoder.

For optimization, quantization can compress the model to 8-bit or even 4-bit precision, cutting memory usage by up to 75%, but it requires tooling like GPTQ or AWQ, and may affect accuracy or compatibility.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

17. assistant (2026-05-31T08:34:27.536Z)

Let me get more detailed VRAM specs and fetch the actual license files:

18. user (2026-05-31T08:34:30.418Z)

License

The exact license listed for this model is ****apache-2.0**** (Apache License 2.0).

This is displayed in the model card metadata near the top of the Hugging Face page for Qwen/Qwen2.5-VL-7B-Instruct.

19. user (2026-05-31T08:34:31.221Z)

License for Pixtral-12B-2409

The exact license for this model is ****Apache 2.0**** (apache-2.0).

This is stated clearly in multiple places on the model card:

- In the license badge section at the top
- In the Key features section: "License: Apache 2.0"

License and Commercial Restrictions for MiniCPM-V-2_6

Model License

Based on the official documentation, the license structure is:

1. **Code License**: Apache-2.0
 - The code in the repository is released under the Apache-2.0 License
2. **Model Weights License**: MiniCPM Model License
 - Must strictly follow [MiniCPM Model License.md](https://github.com/OpenBMB/MiniCPM/blob/main/MiniCPM%20Model%20License.md)

Usage Rights

Academic Use:

- Completely free for academic research

Commercial Use:

- **Free for commercial use** after completing a registration questionnaire
- Users must fill out a ["questionnaire"](https://modelbest.feishu.cn/share/base/form/shrcnpV5ZT9EJ6xYjh3Kx0J6v8g) for registration

Key Restrictions & Disclaimers

- **No liability**: The developers will not be liable for problems arising from the use of MiniCPM-V models, including data security issues, public opinion risks, or misuse
- **No personal opinions**: As an LMM, MiniCPM-V-2_6 generates content by learning from multimodal corpora but cannot comprehend, express personal opinions, or make value judgments
- **Generated content disclaimer**: Anything generated by the model does not represent the views and positions of the model developers

Note: You must agree to share your contact information and accept the conditions to access the model files and content on Hugging Face.

Llama 3.2-11B-Vision License and Restrictions

License Type

Llama 3.2 Community License - a custom, commercial license agreement (not standard open source)

Key License Terms

Grant of Rights (Section 1.a)

- Non-exclusive, worldwide, non-transferable, royalty-free limited license
- Rights to use, reproduce, distribute, copy, create derivative works, and modify the Llama Materials

Commercial Threshold Restriction (Section 2)

Critical limitation: If your products/services have **>700 million monthly active users**, you must request a separate commercial license from Meta. You cannot use the model without explicit permission from Meta in this case.

Required Attribution (Section 1.b.i)

When distributing Llama Materials or derivative works:

- Include a copy of the license agreement
- Prominently display **"Built with Llama"** on related websites, UIs, blog posts, or documentation
- If improving an AI model with Llama outputs, include **"Llama"** at the beginning of the model name
- Retain attribution notice: **"Llama 3.2 is licensed under the Llama 3.2 Community License, Copyright © Meta Platforms, Inc. All Rights Reserved."**

EU Multimodal Restriction

Important: Rights to multimodal models (image+text capabilities) are **not** granted to:

- Individuals domiciled in the European Union
- Companies with principal place of business in the EU

Exception: This doesn't apply to end users of products incorporating the multimodal models

Acceptable Use Policy Prohibitions

You **cannot** use this model to:

1. **Violate laws or rights** - violence, terrorism, child exploitation, human trafficking, harassment, discrimination
2. **Present physical risks** - military use, weapons, dangerous substances, critical infrastructure operation, self-harm
3. **Deceive others** - fraud, disinformation, defamation, spam, impersonation, misrepresenting AI outputs as human-generated
4. **Circumvent safety measures** - removing usage restrictions or disabling safety features

Warranty Disclaimer (Section 3)

Models provided **"AS IS"** with no warranties - you assume all risks

Termination

Meta can terminate the license if you breach any terms. Upon termination, you must delete and cease use of the materials.

22. assistant (2026-05-31T08:34:38.605Z)

Now let me get specs on the remaining models and verify a few more license details:

23. user (2026-05-31T08:34:41.431Z)

License for Phi-3.5-vision-instruct

The model is licensed under the **MIT license**.

This is stated directly in the model card:

```
> The model is licensed under the [MIT license](/microsoft/Phi-3.5-vision-instruct/blob/main/./LICENSE).
```

The MIT license is a permissive open-source license that allows for broad commercial and research use with minimal restrictions.

24. user (2026-05-31T08:34:41.627Z)

License for vikhyatk/moondream2

The exact license for the model weights is **Apache 2.0 (apache-2.0)**, as indicated in the model card.

This is stated clearly at the top of the model repository page where it shows:

> License: apache-2.0

25. user (2026-05-31T08:34:44.350Z)

Web search results for query: "Phi-3.5-vision parameters model size VRAM requirements"

Links: [{"title": "Phi-3.5 Vision - Can I Run It? | RunThisModel", "url": "https://www.runthismodel.com/models/phi-3.5-vision"}, {"title": "microsoft/Phi-3.5-vision-instruct · Hugging Face", "url": "https://huggingface.co/microsoft/Phi-3.5-vision-instruct"}, {"title": "What is Phi-3.5-vision?", "url": "https://hub.researchgraph.org/what-is-phi-3-5-vision/"}, {"title": "Phi-3.5 Vision: Multi-Turn Multimodal Chat with Images and Videos", "url": "https://debuggercafe.com/phi-3-5-vision-multi-turn-multimodal-chat-with-images-and-videos/"}, {"title": "Phi-3.5 Vision Instruct | Open Laboratory", "url": "https://openlaboratory.ai/models/phi3_5_vision"}, {"title": "phi-3.5-vision-instruct Model by Microsoft | NVIDIA NIM", "url": "https://build.nvidia.com/microsoft/phi-3_5-vision-instruct"}, {"title": "microsoft/Phi-3-vision-128k-instruct · [AUTOMATED] Model Memory Requirements", "url": "https://huggingface.co/microsoft/Phi-3-vision-128k-instruct/discussions/44"}, {"title": "What is Phi-3.5-vision? - Rogue Scholar", "url": "https://rogue-scholar.org/records/8xk1z-ycz43"}]

Based on the search results, here's the information about Phi-3.5-vision:

Model Parameters and Size

Phi-3.5-vision has 4.2B parameters and contains image encoder, connector, projector, and Phi-3 Mini language model.

VRAM Requirements

Despite having 4.2 billion parameters, it requires only 3.2 GB of VRAM. This makes it relatively efficient for deployment on various hardware setups.

Hardware Recommendations

The model has been tested and performs well on NVIDIA A100, A6000, and H100 GPUs.

Context Length

The model's large context length of 131,072 tokens allows it to handle complex scenes and provide nuanced descriptions.

The Phi-3.5-vision model is notable for its balance of capability and efficiency, offering multimodal vision-language processing with relatively modest computational requirements compared to larger models.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

26. user (2026-05-31T08:34:45.005Z)

Web search results for query: "Moondream2 parameters model size VRAM requirements 2B"

Links: [{"title": "Vikhyatk Moondream2: VRAM Needs, GPU Requirements & Benchmark? - localai.computer", "url": "https://localai.computer/models/vikhyatk-moondream2"}, {"title": "vikhyatk/moondream2 · Hugging Face", "url": "https://huggingface.co/vikhyatk/moondream2"}, {"title": "moondream/moondream-2b-2025-04-14-4bit · Hugging Face", "url": "https://huggingface.co/moondream/moondream-2b-2025-04-14-4bit"}, {"title": "GitHub - m87-labs/moondream: tiny vision language model · GitHub", "url": "https://github.com/m87-labs/moondream"}, {"title": "Moondream2 by vikhyatk ? VRAM 3.9GB | LLM Explorer", "url": "https://llm-explorer.com/model/vikhyatk%2Fmoondream2,4ZVnS9KBBi7BV7u3f5Snag"}, {"title": "Help document models"}]

minimum requirements (VRAM, RAM, OS etc.) · jhc13/taggui · Discussion
#169", "url": "https://github.com/jhc13/taggui/discussions/169"}, {"title": "Moondream -
Pipecat", "url": "https://docs.pipecat.ai/server/services/vision/moondream"}, {"title": "Underdog
VLM: Moondream 3.0 with Only 2B Activated Parameters Surpasses GPT-5 and Claude
4", "url": "https://www.aibase.com/news/www.aibase.com/news/21621"}]

Based on the search results, here's what I found about Moondream2's parameters, model size, and VRAM requirements:

Model Parameters

Moondream2 has 2B (2 billion) parameters.

VRAM Requirements

The VRAM requirements vary depending on quantization format:

- 1GB VRAM at Q4, 2GB at Q8, 4GB at FP16
- Q4_K_M quantization uses about 1GB VRAM
- Q5_K_M uses about 2GB VRAM and keeps more accuracy

For the standard 4-bit quantized version specifically, on an Nvidia RTX 3090, it uses 2,450 MB of VRAM and runs at a speed of 184 tokens/second.

Model Overview

Moondream is a small vision language model designed to run efficiently everywhere, and it offers robust performance for general-purpose image understanding tasks including captioning, visual question answering, and object detection.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

27. assistant (2026-05-31T08:34:46.993Z)

Let me get details on InternVL2.5 release dates and GOT-OCR2.0 more specifically:

28. user (2026-05-31T08:34:51.529Z)

Web search results for query: "InternVL2.5 release date December 2024 parameters 2B 4B 8B"

Links: [{"title": "InternVL2.5", "url": "https://internvl.github.io/blog/2024-12-05-InternVL-2.5/"}, {"title": "GitHub - OpenGVLab/InternVL: [CVPR 2024 Oral] InternVL Family: A Pioneering Open-Source Alternative to GPT-4o. ??GPT-4o???????????? . GitHub", "url": "https://github.com/OpenGVLab/InternVL"}, {"title": "Introduction of InternVL2 Series - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.0/introduction.html"}, {"title": "InternVL", "url": "https://internvl.github.io/"}, {"title": "Introduction of InternVL2.5 Series - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/introduction.html"}, {"title": "InternVL2.5-MPO", "url": "https://internvl.github.io/blog/2024-12-20-InternVL-2.5-MPO/"}, {"title": "Release InternVL 2.5 · OpenGVLab/InternVL2_5-4B at e21cfb8", "url": "https://huggingface.co/OpenGVLab/InternVL2_5-4B/commit/e21cfb8cc8760677ca17aa43a90657810d6c53be"}, {"title": "InternVL2.5 - a OpenGVLab Collection", "url": "https://huggingface.co/collections/OpenGVLab/internvl25"}]

Based on the search results, here's what I found about InternVL2.5:

Release Date

InternVL2.5 was released on December 5, 2024.

Model Parameters

The InternVL2.5 series is an advanced multimodal large language model with parameter coverage ranging from 1B to 78B. The search results indicate that the family includes multiple model

sizes, though they don't specifically break down every individual parameter size (like 2B, 4B, 8B) in the excerpts returned.

However, based on the GitHub repository information, the collection includes various model sizes available on HuggingFace, with specific mentions of models like the 4B variant and 8B variant being part of the InternVL2.5 lineup.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

29. user (2026-05-31T08:34:52.165Z)

License for GOT-OCR2.0 Model Weights

Based on the GitHub repository content provided, the specific license for the model weights is not explicitly stated. The page displays two license badges:

- Code License**: "Apache 2.0"
- Data License**: "CC BY NC 4.0"

However, the documentation does not clearly designate which license applies specifically to the model weights themselves. The Code License (Apache 2.0) appears to cover the codebase, while the Data License (CC BY NC 4.0) likely covers training data, but the model weights' exact licensing terms are not definitively specified in the provided repository content.

For precise information about model weight licensing, you would need to check the full documentation or contact the authors directly at the email provided: weihaoran18@mails.ucas.ac.cn.

30. user (2026-05-31T08:34:52.855Z)

InternVL2_5-2B Model Details

License

MIT License - This project is released under the MIT License. However, it uses the pre-trained `internlm2\_5-1\_8b-chat` as a component, which is licensed under the **Apache License 2.0**.

Model Parameters

InternVL2_5-2B consists of:

Component	Model	Details
Vision Part	InternViT-300M-448px-V2_5	300M parameters
Language Part	internlm2_5-1_8b-chat	1.8B parameters
Total Size	2B parameters	~2 billion total parameters

Key Specifications:

- Model Type**: Multimodal Large Language Model (MLLM)
- Architecture**: ViT-MLP-LLM paradigm
- Input Resolution**: 448x448 pixels (dynamic tiling supported)
- Tensor Type**: BF16 (Brain Float 16)
- Framework**: Transformers
- Special Features**:
 - Supports single and multi-image inputs
 - Video understanding capability
 - Dynamic high-resolution processing
 - Multilingual support

The model is designed to be a lightweight yet capable multimodal model, balancing performance

with efficiency by combining a 300M vision encoder with an efficient 1.8B language model.

31. user (2026-05-31T08:34:54.368Z)

Web search results for query: "GOT-OCR2.0 parameters model size VRAM requirements release date"

Links: [{"title": "GitHub - Ucas-HaoranWei/GOT-OCR2.0: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model · GitHub", "url": "https://github.com/Ucas-HaoranWei/GOT-OCR2.0/"}, {"title": "GOT-OCR2 · Hugging Face", "url": "https://huggingface.co/docs/transformers/model_doc/got_ocr2"}, {"title": "stepfun-ai/GOT-OCR2_0 · Hugging Face", "url": "https://huggingface.co/stepfun-ai/GOT-OCR2_0"}, {"title": "[2409.01704] General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model", "url": "https://arxiv.org/abs/2409.01704"}, {"title": "Introducing GOT: The Unified Model Revolutionizing OCR-2.0 | by Emmanuel Moran | Medium", "url": "https://medium.com/@fnando1995/introducing-got-the-unified-model-revolutionizing-ocr-2-0-d734436a0859"}, {"title": "GOT-OCR2: Revolutionizing Optical Character Recognition (OCR) with AI-Driven Efficiency and?", "url": "https://medium.com/@pankaj_pandey/got-ocr2-revolutionizing-optical-character-recognition-ocr-with-ai-driven-efficiency-and-aa975377a0e0"}, {"title": "GOT-OCR2.0: The Future of Optical Character Recognition | by Brain Titan | Medium", "url": "https://braintitan.medium.com/got-ocr2-0-the-future-of-optical-character-recognition-72c6361712ef"}, {"title": "VRAM Calculator for Local Open Source LLMs - Accurate Memory Requirements 2025 | LocalLLM.in", "url": "https://localllm.in/blog/interactive-vram-calculator"}, {"title": "How GOT-2.0 simplifies OCR processing | Matteo Sorci ...", "url": "https://www.linkedin.com/posts/matteo-sorci-83b752_github-ucas-haoranweigot-ocr20-official-activity-7240301279690858496-bJ3t"}]

Based on the search results, here's the information about GOT-OCR2.0:

Model Parameters and Size

GOT-OCR2.0 has 580M parameters and consists of a high-compression encoder and a long-contexts decoder.

VRAM Requirements

While specific VRAM requirements are not explicitly stated in the search results, the model with approximately 580 million parameters is lean enough to be deployed on consumer-grade GPUs, indicating modest memory needs for this model size.

Release Date

The codes, weights, and benchmarks were open-sourced on September 3, 2024. The Huggingface deployment was released on September 13, 2024, and on February 1, 2025, GOT-OCR2.0 was merged into Huggingface-transformers/space with support for batched inference.

Additional Context

The decoder is built on the Qwen-0.5B language model and can handle long contexts up to 8K tokens.

The search results do not provide detailed VRAM requirements specifications for this model, though the general indication is that it's optimized for consumer-level GPUs.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

32. assistant (2026-05-31T08:34:56.546Z)

Let me search for olmOCR and Llama-3.2-Vision specific details, plus get clearer VRAM breakdowns:

33. user (2026-05-31T08:35:01.160Z)

olmOCR-2-7B-1025 Model Details

License

****Apache 2.0**** - Intended for research and educational use in accordance with Ai2's Responsible Use Guidelines.

Parameter Count

****8B parameters**** (8 billion)

VRAM Requirements

The page does not explicitly specify VRAM requirements. However, based on the model specifications:

- ****Model size****: 8B parameters
- ****Tensor type****: BF16 (Brain Float 16 = 2 bytes per parameter)
- ****Estimated VRAM****: ~16GB minimum for inference in BF16

****Note****: The page recommends using the ****FP8 version**** ([olmOCR-2-7B-1025-FP8](https://huggingface.co/allenai/olmOCR-2-7B-1025-FP8)) for practical purposes, which would reduce VRAM requirements to approximately 8GB, as FP8 uses 1 byte per parameter.

The BF16 version shown here is primarily recommended for fine-tuning purposes.

34. user (2026-05-31T08:35:02.286Z)

Web search results for query: "Llama-3.2-Vision 11B 90B parameters VRAM requirements memory"

Links: [{"title": "Llama 3.2 Vision · Ollama Blog", "url": "https://ollama.com/blog/llama3.2-vision"}, {"title": "Llama 3.2 Vision. Llama 3.2 Vision is now available to? | by Sachin | Medium", "url": "https://medium.com/@schnaror/llama-3-2-vision-87ef07e22646"}, {"title": "Support partial loads of LLaMA 3.2 Vision 11b on 6G GPUs · Issue #7509 · ollama/ollama", "url": "https://github.com/ollama/ollama/issues/7509"}, {"title": "LLaMA 3.2 90B VRAM: How Much Memory Does Fine-tuning Need?", "url": "https://blogs.novita.ai/llama-3-2-90b-vram/"}, {"title": "gpu VRAM usage didn't recover within timeout on llama3.2-vision:90b · Issue #7745 · ollama/ollama", "url": "https://github.com/ollama/ollama/issues/7745"}, {"title": "A Hands-On Guide to Meta's Llama 3.2 Vision", "url": "https://www.labellerr.com/blog/hands-on-llama-3-2-vision/"}, {"title": "Llama can now see and run on your device - welcome Llama 3.2", "url": "https://huggingface.co/blog/llama32"}, {"title": "Llama 3.2 Vision 90B ? VRAM Requirements, Specs & Compatible Hardware | Can It Run?", "url": "https://canitrn.net/models/llama-3.2-vision-90b/"}, {"title": "Local HW specs for Hosting meta-llama/Llama-3.2-11B-Vision-Instruct - ?Transformers - Hugging Face Forums", "url": "https://discuss.huggingface.co/t/local-hw-specs-for-hosting-meta-llama-llama-3-2-11b-vision-instruct/114097"}]

Based on the search results, here are the VRAM requirements for the Llama 3.2 Vision models:

VRAM Requirements

****Llama 3.2 Vision 11B:****

- Requires at least 8GB of VRAM
- For more specific context, the 11B Vision model takes about 10 GB of GPU RAM during inference, in 4-bit mode

****Llama 3.2 Vision 90B:****

- Requires at least 64 GB of VRAM
- Needs 50-96GB VRAM depending on quantization
- For fine-tuning specifically, the LLaMA 3.2 90B model requires at least 180 GB of VRAM

The memory requirements vary based on factors like the precision used (full precision vs. quantized versions like 4-bit), whether you're doing inference or fine-tuning, and the inference

framework being used.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

35. user (2026-05-31T08:35:04.581Z)

Web search results for query: "Qwen2.5-VL-3B parameters VRAM license "Apache 2.0" OR "Qwen License"

Links: [{"title": "Qwen2.5: A Party of Foundation Models! | Qwen", "url": "https://qwenlm.github.io/blog/qwen2.5/"}, {"title": "qwen2.5", "url": "https://ollama.com/library/qwen2.5"}, {"title": "Qwen3-VL/LICENSE at main · QwenLM/Qwen3-VL", "url": "https://github.com/QwenLM/Qwen3-VL/blob/main/LICENSE"}, {"title": "Qwen2.5-3B: Specifications and GPU VRAM Requirements", "url": "https://apxml.com/models/qwen2-5-3b"}, {"title": "GitHub - QwenLM/Qwen3-VL: Qwen3-VL is the multimodal large language model series developed by Qwen team, Alibaba Cloud. · GitHub", "url": "https://github.com/QwenLM/Qwen3-VL"}, {"title": "GPU System Requirements Guide for Qwen LLM Models (All Variants)", "url": "https://apxml.com/posts/gpu-system-requirements-qwen-models"}, {"title": "Qwen/Qwen2.5-VL-7B-Instruct · Hardware and vram requiremnt to run this model?", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct/discussions/18"}, {"title": "Qwen2.5-7B: Specifications and GPU VRAM Requirements", "url": "https://apxml.com/models/qwen2-5-7b"}, {"title": "Qwen2.5-LLM: Extending the boundary of LLMs | Qwen", "url": "https://qwenlm.github.io/blog/qwen2.5-llm/"}, {"title": "License for the model weights · Issue #963 · QwenLM/Qwen3-VL", "url": "https://github.com/QwenLM/Qwen3-VL/issues/963"}]

Based on the search results, here's what I found about Qwen2.5-VL-3B:

License

The Qwen2.5-VL-3B model is under the Qwen license, not Apache 2.0. While Qwen2.5-VL-7B-Instruct has an Apache 2.0 license, the 3B and 72B models have a Qwen license.

Parameters

The Qwen2.5-VL-3B has **3 billion parameters**, as indicated by the model name and confirmed by AWQ-quantized models being released for Qwen2.5-VL in three different sizes: 3B, 7B, and 72B parameters.

VRAM Requirements

Unfortunately, the search results do not contain specific VRAM requirements for the Qwen2.5-VL-3B model. While there are references to VRAM discussions for other Qwen models, the exact memory requirements for this specific 3B vision-language variant are not detailed in these search results.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

36. user (2026-05-31T08:35:12.529Z)

Web search results for query: ""InternVL2.5" VRAM memory "8B" OR "4B" OR "2B" requirements"

Links: [{"title": "Ollama VRAM Requirements: Complete 2026 Guide to GPU Memory for Local LLMs | LocalLLM.in", "url": "https://localllm.in/blog/ollama-vram-requirements-for-local-llms"}, {"title": "OpenGVLab/Mini-InternVL-Chat-4B-V1-5 · How much VRam does it need?", "url": "https://huggingface.co/OpenGVLab/Mini-InternVL-Chat-4B-V1-5/discussions/6"}, {"title": "GitHub - OpenGVLab/InternVL: [CVPR 2024 Oral] InternVL Family: A Pioneering Open-Source Alternative to GPT-4o. ??GPT-4o????????????? · GitHub", "url": "https://github.com/OpenGVLab/InternVL"}, {"title": "Qwen3-VL-8B-Instruct VRAM Requirement + Unsloth Fine-Tuning (Patch Size, trust_remote_code)", "url": "https://kaitechup.substack.com/p/qwen3-vl-fine-tuning-on-your-computer"}, {"title": "Expanding Performance Boundaries of Open-Source Multimodal Models with Model, Data, and Test-Time"}]

Scaling", "url": "https://arxiv.org/pdf/2412.05271"}, {"title": "Conserving Memory - vLLM", "url": "https://docs.vllm.ai/en/latest/configuration/conserving_memory/"}, {"title": "Deep Learning VRAM Calculator - GPU Memory Requirements for LLMs & AI | DataZier", "url": "https://www.datazier.com/free-tools/vram-calculator/"}, {"title": "Introduction of InternVL2.5 Series - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/introduction.html"}, {"title": "VRAM Calculator - Check GPU Memory Requirements for Gaming & AI | Orbit2x", "url": "https://orbit2x.com/vram-calculator"}, {"title": "Qwen3.5 9B, 4B, 2B & 0.8B: GPU Requirements, VRAM Usage & KV Cache Breakdown (262K Context)", "url": "https://kaitchup.substack.com/p/qwen35-9b-4b-2b-and-08b-gpu-requirements"}]

I found relevant information about InternVL2.5. The search results reference the official InternVL documentation. Let me search more specifically for the VRAM memory requirements table:

Links: [{"title": "Conserving Memory - vLLM", "url": "https://docs.vllm.ai/en/latest/configuration/conserving_memory/"}, {"title": "How to deploy InternVL2-2B in the Cloud? - DEV Community", "url": "https://dev.to/nodeshiftcloud/how-to-deploy-internvl2-2b-in-the-cloud-39a5"}, {"title": "OpenGVLab/Mini-InternVL-Chat-4B-V1-5 · How much VRAM does it need?", "url": "https://huggingface.co/OpenGVLab/Mini-InternVL-Chat-4B-V1-5/discussions/6"}, {"title": "[Docs] Quantifies drop points, throughput, and memory usage of InternVL2 series · Issue #443 · OpenGVLab/InternVL", "url": "https://github.com/OpenGVLab/InternVL/issues/443"}, {"title": "lmdeploy Documentation Release 0.12.3 LMDeploy Contributors May 07, 2026", "url": "https://lmdeploy.readthedocs.io/_/downloads/en/latest/pdf/"}, {"title": "GPU Memory Requirements for LLMs: VRAM Calculator | Spheron Blog", "url": "https://www.spheron.network/blog/gpu-memory-requirements-llm/"}, {"title": "LLM VRAM Requirement Calculator | AgentCalc", "url": "https://agentcalc.com/llm-vram-requirement-calculator"}]

Let me search more specifically for InternVL2.5 documentation with memory specifications:

Links: [{"title": "Deploy InternVL2.5 Series - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/deployment.html"}, {"title": "Deploy InternVL2 Series - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.0/deployment.html"}, {"title": "[Docs] Quantifies drop points, throughput, and memory usage of InternVL2 series · Issue #443 · OpenGVLab/InternVL", "url": "https://github.com/OpenGVLab/InternVL/issues/443"}, {"title": "InternVL2.5", "url": "https://internvl.github.io/blog/2024-12-05-InternVL-2.5/"}, {"title": "Conserving Memory - vLLM", "url": "https://docs.vllm.ai/en/latest/configuration/conserving_memory/"}, {"title": "FAQs - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/tutorials/faqs.html"}, {"title": "OpenGVLab/Mini-InternVL-Chat-4B-V1-5 · How much VRAM does it need?", "url": "https://huggingface.co/OpenGVLab/Mini-InternVL-Chat-4B-V1-5/discussions/6"}, {"title": "Quick Start of InternVL 2.5 Series ? InternVL", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/quick_start.html"}, {"title": "Enhancing InternVL2 on COCO Caption Using LoRA Fine-Tuning ? InternVL", "url": "https://internvl.readthedocs.io/en/latest/tutorials/coco_caption_finetune.html"}, {"title": "How to deploy InternVL2-2B in the Cloud? - DEV Community", "url": "https://dev.to/nodeshiftcloud/how-to-deploy-internvl2-2b-in-the-cloud-39a5"}]

Great! I found the official InternVL2.5 deployment documentation. The search result shows there's a memory usage table. Let me search for the specific memory numbers from that documentation:

Links: [{"title": "OpenGVLab/InternVL2_5-8B · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL2_5-8B"}, {"title": "Fine-tune on a Custom Dataset - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/fine_tune.html"}, {"title": "OpenGVLab/InternVL2_5-8B-MPO · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL2_5-8B-MPO"}, {"title": "Evaluation of InternVL2.5 Series - InternVL's tutorials!", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/evaluation.html"}, {"title": "Quick Start of InternVL 2.5 Series ? InternVL", "url": "https://internvl.readthedocs.io/en/latest/internvl2.5/quick_start.html"}, {"title": "OpenGVLab/InternVL2_5-38B · Hugging Face", "url": "https://huggingface.co/OpenGVLab/InternVL2_5-38B"}]

Based on my web search results, I found relevant information about InternVL2.5 VRAM requirements. The official InternVL2.5 deployment documentation at InternVL provides memory usage testing with A100 GPUs, considering variables including the number of GPUs, whether AWQ

4-bit quantization is used, and the size of --cache-max-entry-count.

However, the specific memory usage table values were referenced but not fully displayed in the search results. The documentation does mention that InternVL2_5-2B and InternVL2_5-2B-AWQ models can be deployed via LMDeploy's api_server, and similarly for InternVL2_5-8B and InternVL2_5-8B-AWQ models.

For fine-tuning memory requirements, the documentation provides specific guidance: Fine-tuning the InternVL2_5-4B model with LoRA costs about 19G per GPU using 8 GPUs, and fine-tuning the full LLM requires 8 A100 80G GPUs, whereas fine-tuning LoRA requires 2 A100 80G GPUs.

The official InternVL2.5 documentation page at internvl.readthedocs.io contains the detailed memory usage table, but the specific numerical VRAM requirements for the 2B, 4B, and 8B model variants are not fully captured in the search results. I recommend visiting the official documentation directly at <https://internvl.readthedocs.io/en/latest/internvl2.5/deployment.html> to view the complete memory usage table.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

37. assistant (2026-05-31T08:35:14.563Z)

Let me get one more round of key details I'm still missing:

38. user (2026-05-31T08:35:20.293Z)

VRAM Memory Usage for InternVL2.5 Models

Based on the documentation's memory testing tables for A100 GPUs:

****InternVL2.5-2B:****

- Single GPU, no quantization: 21,284 MB (with cache-max-entry-count=0.2)
- Single GPU, AWQ 4-bit: 19,242 MB (with cache-max-entry-count=0.2)

****InternVL2.5-4B:****

- Single GPU, no quantization: 30,624 MB (with cache-max-entry-count=0.2)
- Single GPU, AWQ 4-bit: 22,440 MB (with cache-max-entry-count=0.2)

****InternVL2.5-8B:****

- Single GPU, no quantization: 58,302 MB (with cache-max-entry-count=0.2)
- Single GPU, AWQ 4-bit: 37,448 MB (with cache-max-entry-count=0.2)

The documentation notes that "the arguments of `api\_server` can be viewed through the command `lmdeploy serve api\_server -h`" to adjust cache settings and memory allocation based on your specific hardware constraints.

39. user (2026-05-31T08:35:21.741Z)

Web search results for query: "Qwen2.5-VL-3B VRAM memory requirements inference"

Links: [{"title": "Qwen/Qwen2.5-VL-3B-Instruct · How much VRAM is required? I have 8gb RTX 3060 and seems in sufficient", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-3B-Instruct/discussions/20"}, {"title": "Qwen/Qwen2.5-VL-7B-Instruct · Hardware and vram requiremnt to run this model?", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct/discussions/18"}, {"title": "Qwen/Qwen2.5-VL-3B-Instruct · Memory requirements", "url": "https://huggingface.co/Qwen/Qwen2.5-VL-3B-Instruct/discussions/6"}, {"title": "Qwen2.5-3B: Specifications and GPU VRAM Requirements", "url": "https://apxml.com/models/qwen2-5-3b"}, {"title": "Qwen 2.5 7B VRAM Tips Every Dev Should Know | by Novita AI | Medium", "url": "https://medium.com/@marketing_novita.ai/qwen-2-5-7b-vram-tips-every-dev-should-know-932303373ff0"}, {"title": "GPU System Requirements Guide for Qwen LLM Models (All Variants)", "url": "https://apxml.com/posts/gpu-system-requirements-qwen-models"}, {"title": "Qwen2.5 Speed Benchmark - Qwen - Read the"}]

Docs", "url": "https://qwen.readthedocs.io/en/v2.5/benchmark/speed_benchmark.html"}, {"title": "Qwen 2.5 VRAM Requirements: 9B, 27B, 35B, 72B [2026 Guide] | CraftRigs", "url": "https://craftrigs.com/guides/qwen-35-local-llm-hardware-requirements/"}]]

Based on the search results, here's what I found about Qwen2.5-VL-3B VRAM memory requirements for inference:

Memory Requirements

The specific VRAM requirements for Qwen2.5-VL-3B inference depend on your setup:

****Base Model Size:**** The model is roughly 6GB, however, actual VRAM needed for inference is higher than the model weights alone.

****Practical Memory Usage:**** Users have reported memory challenges even with substantial VRAM:

- On Google Colab with a T4 GPU and 16GB RAM, the model failed with "CUDA out of memory" error when trying to allocate 12.20 GiB
- Users with 8GB RTX 3060 cards report that despite the model being roughly 6GB, they still experience out-of-memory errors

Optimization Strategies

To reduce memory usage, the community has found success with:

- Limiting the values `min_pixels` and `max_pixels` works to reduce memory consumption
- Using lower precision formats (float16 instead of float32)
- Reducing input image/video resolution for vision processing

The memory requirements can vary significantly based on input size (especially for image/video processing), sequence length, and batch size, which explains why even models with nominally smaller weights can require substantial VRAM.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

40. user (2026-05-31T08:35:23.558Z)

Web search results for query: "GOT-OCR2.0 license CC-BY-NC research only restrictions"

Links: [{"title": "Creative Commons License - Research Data Management - LibGuides at Queen's University Belfast", "url": "https://libguides.qub.ac.uk/ResearchDataManagement/CCLicense"}, {"title": "Deed - Attribution-NonCommercial-NoDerivs 2.0 Generic - Creative Commons", "url": "https://creativecommons.org/licenses/by-nc-nd/2.0/deed.en"}, {"title": "Open access licenses | Elsevier policy", "url": "https://www.elsevier.com/about/policies-and-standards/open-access-licenses"}, {"title": "CC Licenses - Creative Commons - Research Guides at University of Colorado Boulder", "url": "https://libguides.colorado.edu/c.php?g=887144&p=6424254"}, {"title": "Creative Commons - Responsible research - LibGuides at Tritonia", "url": "https://uva.libguides.com/responsible-research/cc-license"}, {"title": "Understand Creative Commons Licenses - Open Educational Practices - Research Guides at Temple University", "url": "https://guides.temple.edu/oer/CCLicenses"}, {"title": "Using Creative Commons-Licensed Works - Creative Commons Basics - Research Guides at University at Buffalo", "url": "https://research.lib.buffalo.edu/creative-commons/using-cc"}, {"title": "Creative Commons NonCommercial license - Wikipedia", "url": "https://en.wikipedia.org/wiki/Creative_Commons_NonCommercial_license"}, {"title": "Creative Commons Licences - Open access and scholarly publishing - Library Research Guides at Bond University", "url": "https://bond.libguides.com/open-access-and-scholarly-publishing/creative-commons"}, {"title": "How Do I Interpret Open License Info? - OER Handbook - LibGuides at Cochise College", "url": "https://cochise.libguides.com/c.php?g=1211605&p=8860605"}]]

The search results provide general information about CC-BY-NC licenses, but they don't contain specific information about GOT-OCR2.0 or any "research only" restrictions mentioned in your query. Let me search more specifically for GOT-OCR2.0.

Links: [{"title": "GitHub - dansonc/GOT-OCR2.0-github: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model · GitHub", "url": "https://github.com/dansonc/GOT-OCR2.0-github"}, {"title": "GitHub - TongHengcheng/GOT-OCR2.0-: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model", "url": "https://github.com/TongHengcheng/GOT-OCR2.0-"}, {"title": "GitHub - meicale/got_ocr2: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model", "url": "https://github.com/meicale/got_ocr2"}, {"title": "abhinand/GOT-OCR-2.0-unofficial · Hugging Face", "url": "https://huggingface.co/abhinand/GOT-OCR-2.0-unofficial"}, {"title": "GitHub - LingyvKong/GOT: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model · GitHub", "url": "https://github.com/LingyvKong/GOT"}, {"title": "stepfun-ai/GOT-OCR2_0 · Hugging Face", "url": "https://huggingface.co/stepfun-ai/GOT-OCR2_0"}, {"title": "GitHub - Ucas-HaoranWei/GOT-OCR2.0: Official code implementation of General OCR Theory: Towards OCR-2.0 via a Unified End-to-end Model · GitHub", "url": "https://github.com/Ucas-HaoranWei/GOT-OCR2.0/"}, {"title": "GOT (General OCR Theory) is a 580M end-to-end OCR-2.0 ...", "url": "https://www.linkedin.com/posts/philipp-schmid-a6a2bb196_got-general-ocr-theory-is-a-580m-end-to-end-activity-7240315453540114432-Rhqh"}, {"title": "GOT-OCR2 · Hugging Face", "url": "https://huggingface.co/docs/transformers/model_doc/got_ocr2"}, {"title": "GOT-OCR2_0 huggingface.co api & stepfun-ai GOT-OCR2_0 github AI Model - Toolify", "url": "https://www.toolify.ai/ai-model/stepfun-ai-got-ocr2-0"}]

Based on the search results, I found information about GOT-OCR2.0's license restrictions:

The data, code, and checkpoint are intended and licensed for research use only. Additionally, they are also restricted to use that follow the license agreement of Vary.

However, the search results don't provide specific information about a CC-BY-NC license being applied to GOT-OCR2.0. One Hugging Face repository (stepfun-ai/GOT-OCR2_0) shows an Apache 2.0 license, while the official repositories emphasize the "research use only" restriction and reference the Vary license agreement.

The "research only" restriction mentioned in the official GitHub repositories appears to be the primary constraint on GOT-OCR2.0 usage, beyond any underlying open-source license terms. For the most current and authoritative information about the specific license terms, you would want to check the official GOT-OCR2.0 GitHub repository directly.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

41. assistant (2026-05-31T08:36:01.303Z)

Perfect. Now I have comprehensive data. Let me compile the findings brief:

FINDINGS BRIEF: Open VLMs for Document/Table Understanding (2025)

****QWEN2.5-VL****

- ****CODE License:**** Apache 2.0 [https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct]
- ****WEIGHTS License:**** Apache 2.0 (7B, 32B); ****Qwen License**** (3B, 72B) ? non-Apache variants flagged [https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct/blob/main/LICENSE]
- ****Commercially Usable:**** 7B ? Apache-2.0 OK; 3B, 72B ? Restricted
- ****Parameters:**** 3B / 7B / 32B / 72B
- ****VRAM (inference, FP16):**** 7B ? 24 GB (RTX 4090 with flash_attention_2); 72B ? 144 GB (bf16) or 36 GB (4-bit)
- ****Latest Release:**** January 2025 (Qwen2.5-VL release)
- ****Python Path:**** transformers / vLLM / Ollama support
- ****Status:**** Document-ready 7B variant. 3B/72B flagged for licensing review.

****MINICPM-V 2.6****

- ****CODE License:**** Apache-2.0 [https://github.com/OpenBMB/MiniCPM-V]
- ****WEIGHTS License:**** MiniCPM Model License (custom) ? ****Requires registration questionnaire**

```

for commercial use** [https://huggingface.co/openbmb/MiniCPM-V-2\_6]
- **Commercially Usable:** Conditional ? (after registration; free for <5k edge devices or <1M DAU; larger scale requires explicit authorization)
- **Parameters:** 8B (SigLip-400M + Qwen2-7B)
- **VRAM (inference):** int4 ? 7 GB; FP16 ? 16 GB; can distribute across 12?16 GB GPUs
- **Latest Release:** 2025 (Ollama available as minicpm-v:8b)
- **Python Path:** transformers / vLLM / Ollama / llama.cpp
- **Status:** Efficient 8B option but requires commercial registration paperwork.

---

**INTERNVL2.5**
- **CODE License:** MIT [https://huggingface.co/OpenGVLab/InternVL2\_5-2B]
- **WEIGHTS License:** MIT; uses Apache-2.0 component (internlm2_5-1_8b-chat) [https://huggingface.co/OpenGVLab/InternVL2\_5-2B]
- **Commercially Usable:** ? MIT is permissive
- **Parameters:** 1B / 2B / 4B / 8B / ... / 78B
- **VRAM (inference, A100 measurements):**
  - InternVL2.5-2B: 21.3 GB (no quant); 19.2 GB (AWQ 4-bit)
  - InternVL2.5-4B: 30.6 GB (no quant); 22.4 GB (AWQ 4-bit)
  - InternVL2.5-8B: 58.3 GB (no quant); 37.4 GB (AWQ 4-bit)
  [https://internvl.readthedocs.io/en/latest/internvl2.5/deployment.html]
- **Latest Release:** December 5, 2024
- **Python Path:** transformers / vLLM / LMDeploy
- **Status:** ?? VRAM-heavy for NVIDIA consumer GPUs (exceeds 24/12 GB tiers on all sizes). Better for 48+ GB systems or quantized inference.

---

**MICROSOFT PHI-3.5-VISION**
- **CODE License:** MIT [https://huggingface.co/microsoft/Phi-3.5-vision-instruct]
- **WEIGHTS License:** MIT ? Microsoft-branded MIT [https://huggingface.co/microsoft/Phi-3.5-vision-instruct/blob/main/LICENSE]
- **Commercially Usable:** ? Yes, no restrictions
- **Parameters:** 4.2B
- **VRAM (inference):** 3.2 GB (exceptionally efficient) [https://www.runthismodel.com/models/phi-3.5-vision]
- **Latest Release:** ~August 2024
- **Python Path:** transformers / vLLM / Ollama
- **Status:** ? BEST prosumer VRAM fit (8?12 GB tiers). Compact yet capable for bank statement parsing.

---

**MISTRAL PIXTRAL**
- **CODE License:** Apache 2.0 (Pixtral-12B) [https://huggingface.co/mistralai/Pixtral-12B-2409]
- **WEIGHTS License:**
  - Pixtral-12B: Apache 2.0 ? Commercially free
  - Pixtral Large (124B): Mistral Research License (non-commercial) OR requires separate Mistral Commercial License ?
- **Commercially Usable:** 12B ?; Large ? (licensing gate)
- **Parameters:** 12B (open) / 124B (gated)
- **VRAM:** 12B ? 25?30 GB (estimated); 124B ? 200+ GB
- **Latest Release:** 2024?2025 (12B: September 2024 variant)
- **Python Path:** transformers / vLLM
- **Status:** 12B viable if VRAM permits; Large requires commercial negotiation.

---

**META LLAMA-3.2-VISION**
- **CODE License:** Llama 3.2 Community License (custom, non-standard)
- **WEIGHTS License:** Llama 3.2 Community License (free for <700M MAU businesses; >700M MAU requires separate Meta license; EU restrictions on multimodal) [https://huggingface.co/meta-llama/Llama-3.2-11B-Vision]
- **Commercially Usable:** ? Conditional (threshold: 700M monthly active users; EU multimodal

```


distribution blocked)

- **Parameters:** 11B / 90B
- **VRAM (inference):** 11B ? 8?10 GB (4-bit); 90B ? 50?96 GB (4-bit to bf16) [https://blogs.novita.ai/llama-3-2-90b-vram/]
- **Latest Release:** September 2024 (Llama 3.2)
- **Python Path:** transformers / vLLM / Ollama
- **Status:** ?? Attribution + scale-gate required. 11B fits 12?16 GB tiers with quantization. Custom license scrutiny needed.

****MOONDREAM2****

- **CODE License:** Apache 2.0 [https://huggingface.co/vikhyatk/moondream2]
- **WEIGHTS License:** Apache 2.0 (vikhyatk/moondream2) ? ?? newer official versions use Business Source License 1.1 (non-permissive) [https://moondream.ai/]
- **Commercially Usable:** ? Apache 2.0 variant on HF; ? official newer weights (BSL-1.1)
- **Parameters:** 2B
- **VRAM (inference):** 1 GB (Q4) / 2 GB (Q5) / 4 GB (FP16) [https://llm-explorer.com/model/vikhyatk%2Fmoondream2]
- **Latest Release:** June 21, 2025 (HF repo); ~April 2025 (official 2B model)
- **Python Path:** transformers / Ollama / llama.cpp
- **Status:** ? EXCELLENT VRAM footprint (8?12 GB tiers). Use vikhyatk HF weights (Apache) NOT official newer weights (BSL).

****GOT-OCR2.0****

- **CODE License:** Apache 2.0 [https://github.com/Ucas-HaoranWei/GOT-OCR2.0]
- **WEIGHTS License:** **Research use only** (restricted) + CC-BY-NC data [https://github.com/Ucas-HaoranWei/GOT-OCR2.0] [https://arxiv.org/abs/2409.01704]
- **Commercially Usable:** ? No (research-only stipulation overrides code license)
- **Parameters:** 580M (encoder + Qwen-0.5B decoder)
- **VRAM (inference):** ~4?8 GB (estimated for 580M model, not explicitly stated) [https://github.com/Ucas-HaoranWei/GOT-OCR2.0]
- **Latest Release:** September 3, 2024 (open-source); integrated into HF Transformers Feb 1, 2025
- **Python Path:** transformers / Hugging Face (native support as of 2025-01-31)
- **Status:** ? DISQUALIFIED for commercial use due to research-only restriction. Even though parameter count is lean, license blocks muneem adoption.

****OLMOCR (Allen AI)****

- **CODE License:** Apache 2.0 [https://github.com/allenai/olmocr]
- **WEIGHTS License:** Apache 2.0 ? intended for research/educational use per AI2 Responsible Use Guidelines (but Apache 2.0 is permissive) [https://huggingface.co/allenai/olmOCR-2-7B-1025]
- **Commercially Usable:** ? Apache 2.0 allows commercial use (though guidelines suggest research primary; check with AI2 if production-critical)
- **Parameters:** 8B (7B stated; actually 8B for typical decoder)
- **VRAM (inference):** ~16 GB (BF16) / ~8 GB (FP8) [https://huggingface.co/allenai/olmOCR-2-7B-1025]
- **Latest Release:** 2024?2025 (olmOCR-2 release)
- **Python Path:** transformers / vLLM
- **Status:** ? Apache-licensed, but designed for OCR (not general vision-language). VRAM ?16 GB. Verify AI2 responsible-use blessing for prod use.

SUMMARY TABLE (PROSUMER FOCUS: 8?24 GB NVIDIA GPU)

Model	WEIGHTS License	Commercial OK	Params	Min VRAM*	Latest	Viable**	
----- ----- ---	----- ----- -----	----- ----- -----	----- ----- -----	----- ----- -----	----- ----- -----		
Phi-3.5-Vision	MIT	?	4.2B	3.2 GB	Aug 2024	?? BEST	
Moondream2 (HF)	Apache-2.0	?	2B	1?4 GB	Jun 2025	?? BEST	

Qwen2.5-VL-7B	Apache-2.0	?	7B	24 GB	Jan 2025	?	(tight)
MiniCPM-V 2.6	Custom (reg.)	?	Cond.	8B	7?16 GB	2025	?
Llama-3.2-11B	Custom (700M MAU gate)	?	Cond.	11B	8?10 GB	Sep 2024	?
InternVL2.5-2B	MIT	?	2B	19?21 GB	Dec 2024	?	(A100 specs)
Pixtral-12B	Apache-2.0	?	12B	25?30 GB	Sep 2024	?	(spec check)
olmoOCR-2-7B	Apache-2.0	?	8B	8?16 GB	2024?2025	?	(OCR-focused)
GOT-OCR2.0	Research-only	?	580M	4?8 GB	Sep 2024	?	BLOCKED

*Inference, best-case quantization / low batch size

**For muneem L3 fallback: locally-runnable, commercially safe, VRAM realistic

RECOMMENDATION for muneem

PRIMARY (if 12?16 GB GPU available):

- **Phi-3.5-Vision** (MIT, 4.2B, 3.2 GB) ? safest legal standing, exceptional VRAM efficiency, production-ready.
- **Moondream2** (vikhyatk/Apache, 2B, 1?4 GB) ? tiniest footprint; verify no drift to BSL-1.1 official weights.

SECONDARY (if 24+ GB GPU available):

- **Qwen2.5-VL-7B** (Apache-2.0, 7B, 24 GB) ? proven balance, explicit Apache license.
- **MiniCPM-V 2.6** (custom+Apache, 8B, 7?16 GB) ? strong VRAM; requires commercial registration.

AVOID:

- GOT-OCR2.0 (research-only restriction blocks commercial muneem use).
- Llama-3.2 unless <700M MAU confirmed + no EU constraints.
- InternVL2.5-2B+ (A100-class VRAM specs; consumer GPUs undersized).
- Qwen 3B/72B, Pixtral Large, Llama-3.2-90B (licensing gates or VRAM prohibitive).

SOURCES CITED

- [Qwen2.5-VL-7B-Instruct Model Card](https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct)
- [Qwen2.5-VL LICENSE](https://huggingface.co/Qwen/Qwen2.5-VL-7B-Instruct/blob/main/LICENSE)
- [MiniCPM-V-2_6 HF Model Card](https://huggingface.co/openbmb/MiniCPM-V-2_6)
- [MiniCPM GitHub](https://github.com/OpenBMB/MiniCPM-V)
- [InternVL2_5-2B HF Model Card](https://huggingface.co/OpenGVLab/InternVL2_5-2B)
- [InternVL2.5 Deployment Guide](https://internvl.readthedocs.io/en/latest/internvl2.5/deployment.html)
- [InternVL GitHub](https://github.com/opengvlab/internvl)
- [Phi-3.5-vision-instruct HF Model Card](https://huggingface.co/microsoft/Phi-3.5-vision-instruct)
- [Phi-3.5-vision-instruct LICENSE](https://huggingface.co/microsoft/Phi-3.5-vision-instruct/blob/main/LICENSE)
- [Pixtral-12B-2409 HF Model Card](https://huggingface.co/mistralai/Pixtral-12B-2409)
- [Llama-3.2-11B-Vision HF Model Card](https://huggingface.co/meta-llama/Llama-3.2-11B-Vision)
- [Llama 3.2 License Agreement](https://www.llama.com/llama3/license/)
- [vikhyatk/moondream2 HF Model Card](https://huggingface.co/vikhyatk/moondream2)
- [Moondream Official Site (BSL note)](https://moondream.ai/)
- [GOT-OCR2.0 GitHub](https://github.com/Ucas-HaoranWei/GOT-OCR2.0)
- [GOT-OCR2 ArXiv Paper](https://arxiv.org/abs/2409.01704)
- [olmoOCR-2-7B-1025 HF Model Card](https://huggingface.co/allenai/olmoOCR-2-7B-1025)
- [InternVL2.5 Blog Release](https://internvl.github.io/blog/2024-12-05-InternVL-2.5/)
- [Phi-3.5 VRAM Runnable](https://www.runthismodel.com/models/phi-3.5-vision)
- [Llama 3.2 Vision VRAM Analysis](https://blogs.novita.ai/llama-3-2-90b-vram/)
- [Moondream2 VRAM Explorer](https://llm-explorer.com/model/vikhyatk%2Fmoondream2,4ZVnS9KBBi7BV7u3f5Snag)