

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_076beba8-59c\agent-a4149fb471ff46962.jsonl`  
- SHA-256: `fca778d99520f4facbddd4ec9afb1d8305eb8d7dea919dc949a40a71610f7e5`  
- Source modified: `2026-05-31T04:18:57+00:00`  
- Imported at: `2026-07-05T16:48:20+00:00`  
- project: `wf_076beba8-59c`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T04:18:11.549Z)

Source Extractor

Research question: "Identify the best LOCAL (self-hosted, single NVIDIA CUDA GPU) vision / document-AI / OCR models ? and, as a clearly-flagged approval-gated fallback only, paid cloud API services ? for parsing BANK STATEMENT PDFs into structured transaction tables (date, narration/description, withdrawal/debit, deposit/credit, running balance), for a LOCAL-FIRST personal-finance tool ("muneem").

Target documents: Indian bank statements (HDFC, Axis, ICICI, AU Small Finance) ? dense, multi-column, often SEMI-RULED or borderless tables, sometimes multiple accounts per statement; both born-digital (text layer present) and the harder cases where pdfplumber's extract_tables/coordinate clustering is unreliable (e.g. Axis, whose column layout varies between statements).

HARD CONSTRAINTS (a model that violates these is unusable for us):

1. MUST run 100% LOCALLY for statement contents (highly sensitive). NO cloud APIs by default. Paid/cloud is acceptable ONLY as an explicitly-approved fallback ? still report the best cloud options but label them clearly as "cloud/sensitive ? approval-gated, not default".
2. MUST be usable from Python (transformers / vLLM / Ollama / ONNX Runtime / native lib).
3. License MUST be permissive & commercially safe for BOTH the CODE and the MODEL WEIGHTS: MIT, BSD, Apache-2.0, ISC, HPND. AVOID and explicitly FLAG anything GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, or custom "community"/acceptable-use licenses with restrictions. This is critical ? many document-AI models and their training datasets are non-commercial (e.g. LayoutLMv3 weights are CC-BY-NC; Surya and Marker use GPL-or-commercial dual licensing; several VLM weights ship custom restrictive licenses). Verify the WEIGHTS license, not just the repo license.
4. Hardware: a single consumer/prosumer NVIDIA GPU. Give VRAM tiers (~8-12 GB consumer, ~16-24 GB) and which models/quantizations fit each.

Evaluate and compare these LOCAL approaches ? for each give: code license, MODEL-WEIGHTS license (state explicitly), latest release / maintenance status as of 2025-2026, VRAM need + quantization options, Python integration path, and reliability specifically on DENSE FINANCIAL TABLES / multi-column statements:

- Open vision-language models (VLMs) for document/table understanding: Qwen2-VL and Qwen2.5-VL (note per-size license ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2 / 2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision (flag the Llama Community License restrictions), Moondream2, GOT-OCR2.0, olmOCR and other OCR-specialized VLMs.
- OCR-specialist / document-AI pipelines: PaddleOCR + PP-Structure (table recognition), docTR, EasyOCR, Tesseract (already used), Microsoft Table Transformer / TATR (check code AND PubTables-1M dataset/weights license), Donut, Surya (FLAG license), Marker (FLAG license), MinerU / PDF-Extract-Kit, RapidOCR.
- Table-structure recognition specifically (cell/grid reconstruction for borderless tables): TATR, PP-Structure, UniTable, etc.

Also assess: for SEMI-RULED / borderless bank tables, is a prompt-driven VLM (page image -> JSON

transactions) more reliable than a classic OCR + table-structure pipeline? What are the accuracy/hallucination risks of VLMs on financial NUMBERS (transposed digits, dropped/duplicated rows), and how should output be VERIFIED? (muneem already machine-checks parses by balance reconciliation ? opening + sum(deposits) - sum(withdrawals) == closing, plus per-row running-balance walks and cross-checks against the bank's printed summary ? so any model's output can be automatically validated, which should inform the recommendation.)

Finally, briefly cover the best PAID/CLOUD fallbacks (clearly flagged approval-gated for sensitive data): AWS Textract (AnalyzeExpense / Tables), Google Document AI, Azure AI Document Intelligence, Reducto, LlamaParse, Anthropic/OpenAI vision ? note any self-hosting option and data-handling terms.

Deliver a RANKED shortlist of the most promising LOCAL options for muneem (NVIDIA GPU, Python, commercial-license-safe), each with license (code + weights), VRAM tier, maturity, strengths/weaknesses for dense bank-statement tables, and a concrete recommendation: which local model(s) to prototype for muneem's statement parsing and how they'd combine with the existing deterministic parsers + balance-reconciliation verification (e.g. VLM as a fallback/cross-check when deterministic parsing fails to reconcile). Cite sources and call out any license uncertainty explicitly."

Fetch and extract key claims from this source:
URL: https://github.com/QwenLM/Qwen3-VL/issues/963
Title: Qwen VL model-weights license (per-size split) ? official QwenLM GitHub issue
Found via: Model-weights license minefield (commercial-safety, decisive) search

Task

- 1. Use WebFetch to retrieve the page content.
- 2. Assess source quality: primary research/institution? secondary reporting? blog/opinion? forum? unreliable?
- 3. Extract 2-5 FALSIFIABLE claims that bear on the research question. Each claim must:
 - be a concrete, checkable statement (not vague generalities)
 - include a direct quote from the source as support
 - be rated central/supporting/tangential to the research question
- 4. Note publish date if available.

If the fetch fails or the page is irrelevant/paywalled, return claims: [] and sourceQuality: "unreliable".

Structured output only.

2. assistant (2026-05-31T04:18:14.159Z)

I'll fetch the GitHub issue page to extract claims relevant to the research question about Qwen VL model-weights licensing.

3. assistant (2026-05-31T04:18:16.656Z)

I'll fetch the GitHub issue page about Qwen VL model-weights licensing.

4. user (2026-05-31T04:18:24.094Z)

Qwen VL Model Licensing Information

Issue Title: "License for the model weights #963"

Original Question (Posted Mar 17, 2025 by @shaform):
The user noted an inconsistency: "Qwen2.5-VL-7B-Instruct seems to have a Apache 2.0 license, but 3B and 72B ones seem to have a Qwen license" and requested clarification on licensing across model sizes.

Licensing Discrepancy Identified

| Model Variant | License |

```
|---|---|
| Qwen2.5-VL-7B-Instruct | Apache 2.0 |
| Qwen2.5-VL-3B-Instruct | Qwen (custom) |
| Qwen2.5-VL-72B | Qwen (custom) |
```

****Direct Quote from Issue:****

"Is this really intended? Could the authors clarify what's the license for the model weights?"

Response Status

****No official QwenLM team response is present**** in the provided content. The issue remains open with no resolution comments, assignees, or labels. The issue shows 379 total issues in the repository, but no maintainer reply is captured in this excerpt.

5. assistant (2026-05-31T04:18:29.709Z)

The fetch returned useful information but no official maintainer response was captured. Let me verify by checking the Hugging Face model cards directly to confirm the per-size license split, since this is the decisive commercial-safety claim.

6. user (2026-05-31T04:18:35.593Z)

GitHub Issue #963 Analysis

****Issue Status:**** Closed

****Comments:**** No comments are visible in the provided content. The issue shows only the initial description posted by user @shaform on March 17, 2025.

****Key Details from Original Post:****

The user raised concerns about inconsistent licensing across Qwen2.5-VL model variants:

- 7B-Instruct model: Apache 2.0 license
- 3B and 72B variants: "Qwen" license

The original poster asked for clarification on whether this licensing difference was intentional.

****Team Response:**** No confirmation or denial from Qwen team members appears in the visible content provided.

****Additional Elements:**** No reactions, Hugging Face model card links, or subsequent discussion threads are shown in the excerpt supplied.

7. assistant (2026-05-31T04:18:40.667Z)

The two fetches give slightly conflicting status (open vs closed) and no captured maintainer comments. The core content is consistent and reliable, though: the original poster (@shaform, Mar 17 2025) documents the per-size license split for Qwen2.5-VL. This is a user-reported observation on the official QwenLM repo, not an official maintainer confirmation. I have enough to extract the falsifiable claims.

This source is best classified as a ****forum**** (a GitHub issue thread / community Q&A), since the substantive content is a user's observation rather than official documentation or a maintainer ruling. The license-split claim it makes is independently verifiable against Hugging Face model cards and is well-corroborated.

8. user (2026-05-31T04:18:45.807Z)

Structured output provided successfully

9. assistant (2026-05-31T04:18:57.207Z)

I've extracted the key claims from the GitHub issue. Summary of findings:

****Source:**** QwenLM/Qwen3-VL issue #963, "License for the model weights" (posted Mar 17, 2025 by @shaform)

****Source quality: forum**** ? This is a GitHub issue thread. The substantive content is a **user's observation**, not an official maintainer ruling. Both fetch attempts confirmed the original post but captured ****no official QwenLM team response**** (and the two fetches disagreed on open vs. closed status). So the per-size license split here is user-reported, not officially confirmed in-thread ? though it is independently verifiable against the Hugging Face model cards.

****Decisive claim for the research question (commercial safety):****

The Qwen2.5-VL family has a ****per-size license split**** ? Qwen2.5-VL-****7B-Instruct** = Apache-2.0****** (commercial-safe), but ****3B and 72B** = custom "Qwen" license****** (restrictive). This directly confirms the research question's warning that VLM weights licenses must be verified **per size**, not assumed from the repo or family name. For muneem on a single consumer/prosumer GPU, the Apache-2.0 ****7B**** variant is the commercially-safe sweet spot; the 3B (which would otherwise fit smaller VRAM tiers) and the 72B both carry the restrictive license and should be flagged.

****Caveat to carry forward:**** the in-thread license claim is user-asserted; the actual weights license should be confirmed on each Hugging Face model card before relying on it. I flagged this uncertainty in the source-quality rating (forum, not primary).