

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_076beba8-59c\agent-a8130dd0f1780cfe1.jsonl`  
- SHA-256: `384e020f89443e9fa5215bf5ac27938bfa649e9a603f4a35e5a496f7956948be`  
- Source modified: `2026-05-31T04:21:05+00:00`  
- Imported at: `2026-07-05T16:48:20+00:00`  
- project: `wf_076beba8-59c`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T04:18:54.079Z)

Source Extractor

Research question: "Identify the best LOCAL (self-hosted, single NVIDIA CUDA GPU) vision / document-AI / OCR models ? and, as a clearly-flagged approval-gated fallback only, paid cloud API services ? for parsing BANK STATEMENT PDFs into structured transaction tables (date, narration/description, withdrawal/debit, deposit/credit, running balance), for a LOCAL-FIRST personal-finance tool ("muneem").

Target documents: Indian bank statements (HDFC, Axis, ICICI, AU Small Finance) ? dense, multi-column, often SEMI-RULED or borderless tables, sometimes multiple accounts per statement; both born-digital (text layer present) and the harder cases where pdfplumber's extract_tables/coordinate clustering is unreliable (e.g. Axis, whose column layout varies between statements).

HARD CONSTRAINTS (a model that violates these is unusable for us):

1. MUST run 100% LOCALLY for statement contents (highly sensitive). NO cloud APIs by default. Paid/cloud is acceptable ONLY as an explicitly-approved fallback ? still report the best cloud options but label them clearly as "cloud/sensitive ? approval-gated, not default".
2. MUST be usable from Python (transformers / vLLM / Ollama / ONNX Runtime / native lib).
3. License MUST be permissive & commercially safe for BOTH the CODE and the MODEL WEIGHTS: MIT, BSD, Apache-2.0, ISC, HPND. AVOID and explicitly FLAG anything GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, or custom "community"/acceptable-use licenses with restrictions. This is critical ? many document-AI models and their training datasets are non-commercial (e.g. LayoutLMv3 weights are CC-BY-NC; Surya and Marker use GPL-or-commercial dual licensing; several VLM weights ship custom restrictive licenses). Verify the WEIGHTS license, not just the repo license.
4. Hardware: a single consumer/prosumer NVIDIA GPU. Give VRAM tiers (~8-12 GB consumer, ~16-24 GB) and which models/quantizations fit each.

Evaluate and compare these LOCAL approaches ? for each give: code license, MODEL-WEIGHTS license (state explicitly), latest release / maintenance status as of 2025-2026, VRAM need + quantization options, Python integration path, and reliability specifically on DENSE FINANCIAL TABLES / multi-column statements:

- Open vision-language models (VLMs) for document/table understanding: Qwen2-VL and Qwen2.5-VL (note per-size license ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2 / 2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision (flag the Llama Community License restrictions), Moondream2, GOT-OCR2.0, olmOCR and other OCR-specialized VLMs.
- OCR-specialist / document-AI pipelines: PaddleOCR + PP-Structure (table recognition), docTR, EasyOCR, Tesseract (already used), Microsoft Table Transformer / TATR (check code AND PubTables-1M dataset/weights license), Donut, Surya (FLAG license), Marker (FLAG license), MinerU / PDF-Extract-Kit, RapidOCR.
- Table-structure recognition specifically (cell/grid reconstruction for borderless tables): TATR, PP-Structure, UniTable, etc.

Also assess: for SEMI-RULED / borderless bank tables, is a prompt-driven VLM (page image -> JSON

transactions) more reliable than a classic OCR + table-structure pipeline? What are the accuracy/hallucination risks of VLMs on financial NUMBERS (transposed digits, dropped/duplicated rows), and how should output be VERIFIED? (muneem already machine-checks parses by balance reconciliation ? opening + sum(deposits) - sum(withdrawals) == closing, plus per-row running-balance walks and cross-checks against the bank's printed summary ? so any model's output can be automatically validated, which should inform the recommendation.)

Finally, briefly cover the best PAID/CLOUD fallbacks (clearly flagged approval-gated for sensitive data): AWS Textract (AnalyzeExpense / Tables), Google Document AI, Azure AI Document Intelligence, Reducto, LlamaParse, Anthropic/OpenAI vision ? note any self-hosting option and data-handling terms.

Deliver a RANKED shortlist of the most promising LOCAL options for muneem (NVIDIA GPU, Python, commercial-license-safe), each with license (code + weights), VRAM tier, maturity, strengths/weaknesses for dense bank-statement tables, and a concrete recommendation: which local model(s) to prototype for muneem's statement parsing and how they'd combine with the existing deterministic parsers + balance-reconciliation verification (e.g. VLM as a fallback/cross-check when deterministic parsing fails to reconcile). Cite sources and call out any license uncertainty explicitly."

Fetch and extract key claims from this source:

****URL:**** <https://arxiv.org/pdf/2511.14998>

****Title:**** FinCriticalED: A Visual Benchmark for Financial Fact-Level OCR (2025)

****Found via:**** Borderless-table accuracy & financial-number hallucination/verification (contrarian/skeptical) search

Task

1. Use WebFetch to retrieve the page content.
2. Assess source quality: primary research/institution? secondary reporting? blog/opinion? forum? unreliable?
3. Extract 2-5 FALSIFIABLE claims that bear on the research question. Each claim must:
 - be a concrete, checkable statement (not vague generalities)
 - include a direct quote from the source as support
 - be rated central/supporting/tangential to the research question
4. Note publish date if available.

If the fetch fails or the page is irrelevant/paywalled, return claims: [] and sourceQuality: "unreliable".

Structured output only.

2. assistant (2026-05-31T04:19:09.247Z)

I'll fetch the source and extract key claims. Let me start by loading the WebFetch tool schema.

3. assistant (2026-05-31T04:19:14.476Z)

I'll fetch the arxiv PDF and extract claims relevant to the research question about bank statement OCR and financial number hallucination/verification.

4. user (2026-05-31T04:19:28.299Z)

FinCriticalED: Financial Fact-Level OCR Benchmark - Key Findings

1. Benchmark Focus

The benchmark evaluates optical character recognition and document parsing on ****financial documents****, specifically assessing accuracy at the fact (cell/value) level rather than character level. It targets real-world financial tables with challenges like dense formatting and complex layouts.

2. Systems Evaluated

****Vision-Language Models (VLMs):****

- Qwen models (Qwen3, Qwen2.5, Qwen2-VL)
- Claude (Anthropic)
- GPT-4o (OpenAI)
- Gemini (Google)
- DeepSeek models
- InternVL
- LLaMA-based models
- MiniCPM

****Traditional OCR & Parsing Tools:****

- PaddleOCR
- Tesseract
- docTR
- Marker
- GOT-OCR
- Cloud services: AWS Textract

3. Key Quantitative Findings

The document references performance comparisons across models but the compressed PDF stream limits exact metric extraction. References indicate "fact-level accuracy" as the primary metric, with models showing varying performance on financial data extraction tasks.

4. Error Types in Financial OCR

The paper identifies critical issues in financial document processing:

- ****Hallucinations**** in generated content
- ****Transposed digits**** in numerical values
- ****Dropped or duplicated rows**** in tables
- ****Numerical errors**** affecting fact accuracy

5. Challenges in Financial Documents

Financial documents present unique difficulties:

- Dense table structures
- Multi-column layouts
- Borderless or semi-ruled formats
- Complex cell relationships requiring precise alignment

6. VLM vs. Traditional OCR Comparison

The benchmark compares modern vision-language models against traditional OCR pipelines, evaluating which approach better handles financial document complexity.

7. Proposed Metric

The paper proposes ****fact-level error detection/accuracy**** as the evaluation metric, moving beyond character-error-rate to assess whether critical financial facts are correctly extracted.

****Publication:**** arXiv:2511.14998v3 (April 2026)

****Authors:**** Led by Yueru He, Sophia Ananiadou and team

[Binary content (application/pdf, 2.2MB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\9fd4426-f234-4e4a-bfb4-1f2c8255c661\tool-results\webfetch-1780201161953-fwgxf.pdf]

5. assistant (2026-05-31T04:19:31.375Z)

The WebFetch gave a useful overview but the compressed PDF stream limited exact metric extraction. The PDF was saved locally. Let me read the actual PDF to extract verbatim quotes and specific numbers.

6. user (2026-05-31T04:19:32.432Z)

pdftoppm failed: Command 'pdftoppm' not found or is in an unsafe location (current directory)

7. assistant (2026-05-31T04:19:34.836Z)

PDF rendering needs poppler. Let me check if I can extract text another way using Python (pypdf is already a dependency of this project per the git status).

8. user (2026-05-31T04:19:36.270Z)

Exit code 1

Traceback (most recent call last):

```
File "<string>", line 1, in <module>
  import pypdf; print(pypdf.__version__)
^^^^^^^^^^^^^^^^
```

ModuleNotFoundError: No module named 'pypdf'

Traceback (most recent call last):

```
File "<string>", line 1, in <module>
  import pdfplumber; print('pdfplumber', pdfplumber.__version__)
^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

ModuleNotFoundError: No module named 'pdfplumber'

9. assistant (2026-05-31T04:19:38.641Z)

Not in the base environment. This project uses a virtualenv likely. Let me check for a venv and try extracting text from the PDF.

10. user (2026-05-31T04:19:41.177Z)

C:/Users/avidu/Projects/muneem/.venv/Scripts/python.exe

===

ok 6.12.2 0.11.9

11. assistant (2026-05-31T04:19:43.007Z)

The project venv has pdfplumber and pypdf. Let me extract the text from the saved PDF.

12. user (2026-05-31T04:19:45.897Z)

NUM PAGES: 13

FinCriticalED: A Visual Benchmark for Financial Fact-Level OCR

YueruHe YupengCao YanWang

ColumbiaUniversity StevensInstituteofTechnology TheFinAI

USA USA USA

yh3507@columbia.edu

LingfeiQian ShuyaoWang YiHan

TheFinAI TheFinAI GeorgiaInstituteofTechnology

USA USA USA

HaohangLi RuoyuXiang FanZhang

StevensInstituteofTechnology NewYorkUniversity TheUniversityofTokyo&MBZUAI

USA USA Japan&UAE

ZhuohanXie MingquanLin PrayagTiwari

MBZUAI UniversityofMinnesota HalmstadUniversity

UAE USA USA

XueqingPeng? JiminHuang GuojunXiong

TheFinAI TheFinAI HarvardUniversity

USA USA USA

xueqing.peng2024@gmail.com

SophiaAnaniadou

UniversityofManchester

UK

Figure1:OverviewofFinCriticalED.Left:standardOCRmetricsfailtocapturefinanciallycriticalsemantic errors.Middle-left:

constructionofafact-centricdatasetwithfivetypesofannotatedfinancialfacts.Middle-

right:evaluationpipelineusing
 OCR/VLM/MLLMoutputswithDeterministic-Rule-GuidedLLM-as-
 Judge.Right:Keyinsightsfromexperimentson13OCR
 systemsandMLLMs.
 Abstract
 whetherOCRandvision-languagepreservethefinancially
 Recentprogressinmultimodal largelanguagemodels(MLLMs)has critical evidence beyond lexical
 similarity. FinCriticalED con-
 substantiallyimproveddocumentunderstanding, yetstrongoptical tains859real-
 worldfinancialdocumentpageswith9,481expert-
 characterrecognition(OCR)performanceonsurfacemetricsdoes
 annotatedfactsspanningfivecriticalfieldtypes:numeric,temporal,
 notguaranteefaithfulpreservationofdecision-criticalevidence.
 monetaryunit,reportingentity,andfinancialconcept.Weformu-
 Thislimitationisespeciallyconsequentialinfinancialdocuments, latethetaskasstructuredOCRwithfact-
 levelverification,and
 wheresmallvisualerrorsscaninducediscreteshiftsinmeaning.To developaDeterministic-Rule-GuidedLLM-
 as-Judgeprotocoltoas-
 studythisgap,weintroduceFinCriticalED(FinancialCritical
 sesswhethermodeloutputspreserveannotatedfactsincontext.We
 ErrorDetection),afact-centricvisualbenchmarkforevaluating
 benchmark13systemsspanningOCRpipelines,specializedOCR
 VLMs,open-sourceMLLMs,andproprietaryMLLMs.Resultsreveal
 ?Correspondingauthor.
 6202
 rpA
 7
]VC.sc[
 3v89941.1152:viXra

13. assistant (2026-05-31T04:19:49.271Z)

I have the abstract. The text extraction is working (though words run together due to font encoding). Let me get the key results pages ? the abstract continues, plus the experiments/results tables. Let me grab the rest of page 1 and search for the quantitative findings.

14. user (2026-05-31T04:19:51.720Z)

(Bash completed with no output)

15. user (2026-05-31T04:19:56.236Z)

LEN P1: 2097
 ''

16. assistant (2026-05-31T04:19:58.951Z)

Page 1 was only 2097 chars (the abstract was cut at "benchmark13systems..."). The rest is on subsequent pages. Let me pull pages 2 and the results section. Let me search across all pages for the key result terms and dump the most relevant pages.

17. user (2026-05-31T04:20:05.711Z)

```
0 len 2097 hits: []
1 len 5666 hits: ['marker', 'f1', 'accuracy']
2 len 5708 hits: ['gemini', 'claude', 'paddleocr', 'textextract', 'deepseek']
3 len 5709 hits: ['gemini', 'gpt-4o', 'claude', 'qwen', 'textextract', 'accuracy', 'deepseek']
4 len 5613 hits: ['gemini', 'gpt-4o', 'claude', 'qwen', 'marker', 'accuracy', 'best',
'outperform', 'deepseek']
5 len 5209 hits: ['gpt-4o', 'marker', 'hallucinat', 'deepseek']
6 len 9342 hits: ['gemini', 'gpt-4o', 'claude', 'qwen', 'paddleocr', 'deepseek']
7 len 5991 hits: ['marker']
8 len 4640 hits: ['accuracy']
```

9 len 6417 hits: ['gpt-4o', 'accuracy']
10 len 5256 hits: ['gpt-4o', 'paddleocr', 'accuracy', 'hallucinat']
11 len 4196 hits: ['gpt-4o', 'claude', 'qwen']
12 len 4093 hits: ['gpt-4o', 'accuracy', 'hallucinat', 'deepseek']

18. assistant (2026-05-31T04:20:08.496Z)

The richest results pages are 4 (best/outperform), 5 and 10/12 (hallucination), 3 and 6 (model lists/tables). Let me extract these key pages.

19. user (2026-05-31T04:20:11.327Z)

===== PAGE 4 =====

FinCriticalED:AVisualBenchmarkforFinancialFact-LevelOCR ?, ,

Table5:ModelperformancesfortraditionalOCRpipelines, specializedOCRVLMS, open-sourceMLLMs, andproprietary MLLMs. Boldindicatesthebestperformanceineachcolumn, whileunderlinedvaluesdenotethesecond-best. Forerror rate(E?), lowerisbetter. Fact-levelmetricsincludenumerical (N-FFA), temporal(T-FFA), monetaryunit(MU-FFA), report-ingentity(RE-FFA), andfinancialconcept(FC-FFA)accuracy; FFAdenotesoverallfinancialfactaccuracy.

Model Size General(%) Fact-Level(%)

ROUGE-1? ROUGE-L? EditDistance? N-FFA? T-FFA? MU-FFA? RE-FFA? FC-FFA? FFA?

OCRPipelines

MinerU2.5[23] 1.2B 95.71 95.30 6.02 98.76 96.48 54.05 91.09 96.44 94.64

PP-OCrv5[4] 0.07B 97.54 96.55 3.10 95.70 90.29 90.00 86.62 93.75 91.91

SpecializedOCRVLMS

DeepSeek-OCR[35] 6B 94.73 94.42 7.33 93.47 91.96 83.53 92.27 94.36 92.67

DeepSeek-OCR-2[36] 3B 92.90 92.18 10.72 82.63 91.90 82.83 88.69 86.51 86.19

GLM-OCR[5] 0.9B 95.10 94.74 6.43 93.24 98.53 88.89 97.84 100.00 96.92

Open-sourceMLLMs

Gemma-3n-E4B-it[10] 4B 83.49 79.59 23.82 52.65 77.06 64.71 74.65 72.86 65.68

Qwen3-VL-8B-Instruct[29] 8B 97.68 97.40 2.93 98.47 96.99 97.65 93.18 99.24 96.88

Llama-4-Maverick[21] 17B 98.00 97.62 3.70 97.77 97.99 97.65 94.26 98.48 96.92

Qwen3.5-397B-A17B[28] 397B 98.12 98.00 2.59 87.72 87.99 86.14 91.22 94.40 89.70

ProprietaryMLLMs

GPT-4o[12] - 90.40 88.35 16.01 59.56 84.59 81.92 81.78 70.84 71.68

GPT-5[24] - 91.81 89.56 15.79 66.83 94.48 92.35 89.19 91.77 81.65

Claude-Sonnet-4.6[1] - 98.84 98.73 1.69 98.59 97.99 97.06 94.02 98.94 97.23

Figure2:ModelperformanceongeneralOCRmetricsand

Gemini-2.5-Pro[9] - 98.81 98.37 2.46 97.24 97.82 97.65 94.18 98.94 96.74

financial fact-level accuracies. Best results per metric are highlightedwithastar(?).

4.1 MainResultsandDiscussion

4.1.1 GeneralOCRMetricsvs. FinancialFidelity. RQ1:To what importantadvantageforfinanciallyfaithfultranscription, although extentdoconventionalOCRmetricsreflectfinanciallyfaithful ex- thestrongestgeneralMLLMscannowmatchorevenexceedit. tractionofcriticalfields?

4.1.3 DifficultyAcrossFinancialCriticalFactTypes. RQ3:Which TheresultsshowthatconventionalOCRmetricsareinsufficient typesoffinancialcriticalfactsaremostvulnerabletoOCRerrors forevaluatingfinancialfidelity. Whiletheycaptureoverallrecon- acrosscurrentmodels?

structionquality, theydonotreliablyreflectwhetherfinanciallycrit- Theresultsshowthatnumericalvaluesandmonetaryunitsare icalfieldsarepreservedfaithfully. Forexample, MinerU2.5 achieves themostvulnerablefacttypesacrosscurrentmodels, whiletem- anear-ceilingROUGE-1scoreof95.71%andROUGE-Lscoreof poral facts and financial concepts perform relatively better but 95.30%, yetdropsharplyto54.05%onMU-FFA, indicatingthat remainfarfromrobust. Forexample, Gemma-3n-E4B-it achieves stronglexicaloverlapcanstillmaskseriousfact-levelerrors. Similar only52.65%, GPT-4o59.56%, andGPT-566.83%onnumericalfacts,

gaps remain even among the strongest models: Claude-Sonnet-4.6 indicating a substantial gap in preserving reliable numeric accuracy? an essential requirement for financial document understanding. MinerU2.5 reaches only 54.05% on MU-FFA, suggesting that and 98.00% ROUGE-L yet falls to 86.14% on MU-FFA. These results nearly half of monetary unit fields are incorrectly recognized, which show that high overlap-based OCR scores do not necessarily imply can severely distort financial meaning. Even Claude-Sonnet-4.6, financially faithful extraction, especially when correctness depends despite its more balanced overall performance, still falls short of on exact values, units, reporting subjects, and field associations. the near-perfect accuracy required for real-world financial applications. 4.1.2 Specialized OCR Models vs. General MLLMs. RQ2: How do cations. Notably, even the comparatively stronger fact categories specialized OCR models and general-purpose MLLMs differ in their fail to reach trustworthy levels. For instance, Qwen3.5-397B-A17B ability to preserve financially critical information? achieves 87.99% on T-FFA and 94.40% on FC-FFA, while weaker In our benchmark, OCR-specialized models show strong performance but not uniformly dominant on financially critical extraction. 77.06% on T-FFA. Overall, these results suggest that compact local Among the OCR-specialized systems, GLM-OCR achieves 96.92% cues, such as digits, decimal points, signs, and unit markers, are FFA, outperforming PP-OCRv5 at 91.91% and DeepSeek-OCR-2 at the most error-sensitive elements in financial OCR, and that even 86.19%, which indicates that OCR-oriented models remain comparatively stronger fact types have not yet achieved the level of petitive baselines for exact field preservation. At the same time, reliability required for high-stakes financial interpretation. several general MLLMs achieve comparable or stronger fact-level 4.2 Model Failure Analysis performance, including Qwen3-VL-8B-Instruct at 96.88%, Llama-4-Maverick at 96.92%, Gemini-2.5-Pro at 96.74%. Among these models, Beyond aggregate fact-level accuracy, we conduct a page-level Claude-Sonnet-4.6 gets the best result with 97.23%. It shows the potential of MLLMs in OCR tasks. However, this capability is reviewed by financial experts. While FFA evaluates whether a not consistent across the broader general-MLLM group: Gemma-3n-notated financial critical facts are preserved, expert review further E4B-it reaches only 65.68% FFA, GPT-4o 71.68%, and GPT-5 81.65%, captures non-critical OCR errors, layout distortions, and other page-level failures that still affect practical use. This means that a model Overall, these findings suggest that OCR specialization remains an may achieve high FFA or ROUGE scores yet still exhibits substantial

===== PAGE 5 =====

?, Heetal.

Table 6: Page-level error statistics across models. Critical error: if at least one annotated financially critical field is incorrect; non-critical error: if OCR discrepancies occur outside annotated critical spans.

Model CriticalError Non-CriticalError CleanPages CriticalOnly Non-Critical Both Rate(%)? Rate(%)? (%)? (%)? Only(%)? (%)?

OCR Pipeline

MinerU2.5[23] 23.5 58.8 32.4 8.8 44.1 14.7

Specialized OCR VLMs

DeepSeek-OCR[35] 17.6 86.3 13.7 0.0 68.6 17.7

Open-source MLLMs

Gemma-3n-E4B-it[10] 74.5 90.2 7.8 2.0 17.7 72.6

Llama-4-Maverick[21] 15.7 49.0 49.0 2.0 35.3 13.7

Proprietary MLLMs

GPT-4o[12] 52.9 80.4 11.8 7.8 35.3 45.1

GPT-5[24] 58.0 86.0 12.0 2.0 30.0 56.0

page-level errors. We analyze these failures along three axes: (i) financial severity, distinguishing between critical-field and non-critical-field errors; (ii) document modality and complexity, examining how performance varies across page types and difficulty levels; and (iii) model-specific failure patterns, characterizing recurring error behaviors across model families. Each subtable reports a 3x3 breakdown examining how performance varies across page types and difficulty levels; and (iii) model-specific failure patterns, characterizing recurring error behaviors across model families. Grey cell with "-" represents 0 sample under this category.

Financial severity. We define a critical error as any OCR mismatch involving one of the five annotated critical field types, and a non-critical error as any OCR discrepancy outside these punctuation or dashes. This is consistent with their architectures: fields, such as errors in surrounding narrative text, punctuation, MinerU2.5 uses a coarse-to-fine two-stage pipeline that separates

or spacing. Table 6 shows large differences in page-level financial layout analysis from local recognition [23], while DeepSeek-OCR reliability: the strongest systems make critical errors on only 17.6% of pages, while weaker baselines do so on nearly three-fourths of pages, while weaker baselines do so on nearly three-fourths of pages with compact visual tokens [35]. By contrast, weaker or more general models produce more diverse and less predictable errors. In Gemma-3n-E4B-it and GPT-4o, we observe hallucinated fact preservation does not imply clean transcription.

continuation, where the model incorrectly completes a sentence. A decomposition into clean, non-critical only, critical only, and using previously seen text with a similar prefix, along with context-both further highlight this gap. Llama-4-Maverick yields the high-free-form corruption [10, 12].

estimated share of clean pages (49.0%), whereas Gemma-3n-E4B-it has the largest share of pages containing both critical and non-critical errors (72.6%). Thus, practical financial OCR requires minimizing both transcription and semantic reconstruction.

materially dangerous errors and broader transcription artifacts.

5 Conclusion

Document modality and complexity. We next stratify page-level critical error rates by document modality and visual complexity. Pages are grouped into text-only, table-only, and mixed modality, serving financially critical evidence in real-world documents. The benchmark includes an expert-annotated dataset of financially critical fields and an LLM-powered evaluation protocol for structured fact-level assessment. Our experiments yield three main findings: setting. Weaker general-purpose LLMs exhibit very high critical error rates in this category, with Gemma-3n-E4B-it reaching financial faithfulness: high surface-level overlap can still mask substantial fact-level errors. (2) Not all critical fields are equally difficult: at 80.0%±10.0%. By contrast, OCR-specialized models are more numerical fields and monetary units are consistently more fragile, though not error-free: on text-only pages, DeepSeek-OCR is less than temporal, reporting, and concept fields, showing that the

ical fields and an LLM-powered evaluation protocol for structured

As shown in Figure 3, mixed pages are consistently the hardest

fact-level assessment. Our experiments yield three main findings:

setting. Weaker general-purpose LLMs exhibit very high critical

(1) Conventional text similarity metrics do not reliably capture error rates in this category, with Gemma-3n-E4B-it reaching

financial faithfulness: high surface-level overlap can still mask sub-

100.0% across all complexity levels and GPT models remaining

stantial fact-level errors. (2) Not all critical fields are equally difficult:

at 80.0%±10.0%. By contrast, OCR-specialized models are more

numerical fields and monetary units are consistently more frag-

stable, though not error-free: on text-only pages, DeepSeek-OCR

is less than temporal, reporting, and concept fields, showing that the

risers from 0.0% at low complexity to 23.1% at high complexity, and main challenge lies in preserving small local cues such as digits, MinerU 2.5 rises from 0.0% to 38.5%. Table-only pages are generally signs, and unit markers. (3) OCR-specialized systems remain strong easier when layout is simple. Overall, the main challenge is OCR and stable baselines, but the strongest LLMs can now match under heterogeneous structure rather than OCR alone. ore exceed them, suggesting that grounded transcription behavior Model-specific failure patterns. Beyond aggregate error rates, matters more than model family alone. Overall, FinCritical ED model families exhibit distinct failure modes. OCR-specialized sys- re frames financial OCR as a fact- preservation problem rather than temss such as DeepSeek-OCR and MinerU 2.5 mostly make fixed, as a surface transcription task. We hope the benchmark will support repetitive recognition errors, including recurring lexical corruptions future research on financially faithful, layout-aware document un- (e.g., transfer or? trans fessor) and stable formatting artifacts around derstanding in precision- critical domains.

20. assistant (2026-05-31T04:20:14.734Z)

This is excellent ? I now have the full results table and the key narrative findings. This is a rich, directly-relevant source. Let me grab the model lists/setup (page 3) and the hallucination discussion (pages 10, 12) to confirm specifics and the LLM-as-Judge / verification angle.

21. user (2026-05-31T04:20:20.278Z)

===== PAGE 3 =====

?, , Heetal.

Table 4: Statistics and financial critical facts distribution of and (3) any entity passing an normalized exact-match check is deter- the FinCritical ED dataset. ministically forced to correct=true, guarding against LLM false Type Count Average per Document Percentage (%) negatives. Rather than penalizing superficial formatting differences, Dataset Size 859 ? ? the judge focuses on semantic fidelity: it determines whether the Numeric Facts(??) 2,930 3.4 30.9 model has successfully located and reproduced each entity, and Temporal Facts(??) 1,144 1.3 12.1 whether there reproduced value is factually exact. Conditioned on a Monetary Units(???) 881 1.0 9.3 structured evaluation prompt (see in Appendix), the Deterministic- Reporting Entities(???) 3,142 3.7 33.1

Financial Concepts(? ??) 1,384 1.6 14.6 Rule-Guided LLM-as-Judge outputs per- field correctness decisions

Total(?) 9,481 11.0 100.0 used to compute accuracy metrics across critical field types.

To rigorously validate the automated evaluator, we conducted a

human? LLM alignment validation by sampling 15 model outputs

3.2 Financial Fact Evaluation

each from PP-OCRv5, GPT-4o, and GPT-5. We adopt a stringent

Traditional OCR evaluation relies on surface-level lexical metrics Document-

Level Strict Agreement metric: the LLM judge and human

(e.g., ROUGE, edit distance) that rewards superficial string overlap.

experts are considered in agreement if and only if their correctness

However, these metrics fail to assess whether a model preserves the decisions(? ?) match perfectly across all entities within a document.

financial meaning of a document. In finance, even single-character Even as a single entity-

level divergence results in a document-level

OCR deviations can alter the underlying reality. FinCritical ED

disagreement. Following iterative prompt and workflow refinement,

operationalizes a shift from text similarity to fact-level correctness,

the evaluation protocol achieves a robust 95% agreement rate, con-

measuring the exact preservation of domain-critical factual content. firming its reliability as a domain-

aware judge for complex financial

factual verification.

3.2.1 Financial Fact Accuracy (FFA) Formulation. The objective is

to verify whether a model correctly reproduces the annotated facts

3.3 Experimental Setup

? =? ??? ??? ???? ???? ?? from the ground-truth HTML within

the model-generated output. For each fact f , we assign a binary value $y_f \in \{0, 1\}$, where $y_f = 1$ if f is recovered comprehensively across all categories in a semantically equivalent local context, and 0 otherwise. Crucially, factual matching across all categories requires the absolute OCR pipelines. Modular systems that decouple text detection, preservation of visually subtle but determinant cues, including recognition, and document reconstruction, including MinerU2.5 [23] signs, decimals, units, scales, temporal boundaries, and entity attributes and PP-OCRv5 [4]. These serve as strong engineering-oriented baselines. For example, converting $(1, 200)$ to $1,200$ yields $y_f = 0$, as lines for structured document parsing. The omission of the parenthetical negative notation fundamentally specializes OCR VLMs. Specialized foundation models optimized to invert the financial meaning for visually grounded text extraction and structure-aware recognition. For a given output document, the overall Financial Fact Annotation, including DeepSeek-OCR [35], DeepSeek-OCR-2 [36], and curacy (FFA) is computed as $\text{curacy}(\text{FFA}) = \frac{1}{N} \sum_{f \in \mathcal{F}} y_f$. We similarly define GLM-OCR [5]. These provide a tightly integrated OCR-central category-specific accuracies as $\text{acc}_c = \frac{1}{N_c} \sum_{f \in \mathcal{F}_c} y_f$ for each subset $\mathcal{F}_c$ of facts. To obtain robust dataset-level performance, we open-source MLLMs. Unified multimodal models that jointly normalized across documents with varying fact densities, we aggregate the total correct counts $\sum y_f$ and total fact counts N across Gemma-3n-E4B-it [10], Qwen3-VL-8B-Instruct [29], Llama-4-Mav-allsamples. The global benchmark metric for each category is thus $\text{erick} = \frac{\sum y_f}{N}$ [21], and Qwen3.5-397B-A17B [28]. Defined as $\text{FFA} = \frac{\sum y_f}{N}$, yielding distinct evaluations for numeric Proprietary MLLMs. Frontier closed-source models with strong (N-FFA), temporal (T-FFA), monetary unit (MU-FFA), reporting multimodal reasoning capabilities, including GPT-4o [12], GPT-entity (RE-FFA), and financial concept (FC-FFA) preservation. The 5 [24], Claude-Sonnet-4.6 [1], and Gemini-2.5-Pro [9]. These serve as overall dataset-level FFA is computed analogously over all facts in upper-bound references for end-to-end document understanding, providing a unified measure of factual fidelity without external OCR.

3.2.2 Deterministic-Rule-Guided LLM-as-Judge Verification.

Because

3.3.2 Implementation Details.

All proprietary models are accessed via their official APIs with temperature set to zero for deterministic outputs. Open-source MLLMs are hosted by Together AI API. fact may admit multiple valid expressions (e.g., $\$1.2\text{B}$ vs. $1,200$ million), whereas a lexically minor deviation can entirely falsify on four NVIDIA GPUs or accessed via their public API endpoints. the preserved information (e.g., $16,950$ vs. 16.950). To evaluate all models share a unified input/output format aligned with the fact level, we employ a Deterministic-Rule-Guided LLM-as-FinCritical ED annotations schema. Judge pipeline combining deterministic parsing with LLM-powered semantic judgment across the five fact types. The pipeline proceeds

4 Results and Analysis

in three stages: (1) Extract all annotated entities from ground-truth HTML with ± 40 -character context hints for disambiguation; (2) To assess whether current OCR and MLLM systems truly preserve LLM-as-Judge framework (powered by GPT-4o) determines per financially critical, visually grounded evidence, we organize our entity whether it was found and exactly matched in the prediction; experiments around the following three research questions.

===== PAGE 10 =====

FinCritical ED: A Visual Benchmark for Financial Fact-Level OCR, judge and human must agree on every entity within a document to count as aligned; a single divergence constitutes disagreement. GPT-

*total_entities counts MUST match the counts from the pre-extracted list.

*REMINDER: if matched_text equals value(normalized), correct MUST be true?

4o achieved 95% agreement after iterative prompt and evaluation

do not set correct=false just because the surrounding context in the prediction differs from context_hint. workflow_refinement.Figure 5 shows a representative alignment

""" example:

Case: Overall FFA=90%. On Figure 5, both human experts and the LLM-as-Judge output sample

LLM-as-Judge classify the output as factually unreliable. The LLM-as-Judge hallucinates the row name "Certificate rate" to "Interest

```
{'total_entities': 65, rate, "as it associates with the "Interest paid" row name below. The
'total_entities_with_Number_type': 37,
'total_entities_with_Temporal_type': 10, model also misreads the monetary amount. LLM output:
'total_entities_with_Monetary_Unit_type': 7,
'total_entities_with_Reporting_Entity_type': 11,
'total_entities_with_Financial_Concepts_type': 0, {
'found_entities': 64, "total_entities": 23,
'found_entities_with_Number_type': 36, ...
'found_entities_with_Temporal_type': 10, "found_entities_with_Number_type": 3,
'found_entities_with_Monetary_Unit_type': 7, ...
'found_entities_with_Reporting_Entity_type': 11,
"found_entities_with_Financial_Concepts_type": 10,
'found_entities_with_Financial_Concepts_type': 0, ...
'correct_entities': 64, "correct_entities_with_Number_type": 2,
'correct_entities_with_Number_type': 36, ...
'correct_entities_with_Temporal_type': 10, "correct_entities_with_Financial_Concepts_type": 9,
'correct_entities_with_Monetary_Unit_type': 7, "entity_accuracy": 90.0,
'correct_entities_with_Reporting_Entity_type': 11,
"overall_explanation": "The model accurately identified most entities,
'correct_entities_with_Financial_Concepts_type': 0,
especially reporting entities and financial concepts, but struggled
'entity_accuracy': 98.36, with exact numeric matches."
'overall_explanation': "The model successfully identified and correctly }
matched nearly all entities, with only one minor discrepancy in a number
format.",
'entity_diagnostics': [{'type': 'Number', Human expert explanation:
'value': '10%',
'context_hint': 'X Director 10% Owner Officer (give title below) Other',
'found_in_prediction': True, Total entities: 23
'correct': True, Numeric entities: 3
... ..
'not_found_entities_with_Number_type': 1, Financial Concepts: 10
'not_found_entities_with_Temporal_type': 0, ...
'not_found_entities_with_Monetary_Unit_type': 0, Wrong numeric entities: 1
'not_found_entities_with_Reporting_Entity_type': 0, ...
'not_found_entities_with_Financial_Concepts_type': 0} Wrong reporting entities: 1
Overall: The model performs well overall, except for errors due to hallucination
on "Interest rate" and wrong numerical entity of "21,829,000
(should be "21,429,000").
Design Details. Context Window. Many financial documents re-
peat identical values across different line items (e.g., the number
These observations correspond precisely to the errors and overall
?1,200? may appear in both revenue and headcount rows of the judgment from GPT-4o output.
same table). To prevent the judge from conflating these, each entity
is accompanied by the 40 characters immediately preceding and D Model Failure Analysis
following it in the gold text. For instance, the entity 1,200 might be
This section analyzes model failures beyond overall FFA scores. We
disambiguate by the hint?...total revenue (in millions)
first describe the error-analysis setup, then examine two notable
1,200 for the fiscal year...?, directing the judge to verify
aggregate patterns: elevated missing-output behavior in some pro-
that is specific occurrence rather than any surface match.
proprietary models and GLM-OCR? s strong performance on financial
Normalization and Deterministic Override. Before applying the
concepts. We finally present representative error cases illustrating
```

override, both the predicted and gold entity values are normal-howdifferentfailuremodesappearinpractice.

ized:lowercased,whitespace-collapsed,andstrippedofpunctua-
tion(commas,currencysymbols,trailingperiods).Ifthenormalized D.1 ErrorAnalysisSetup
stringsareidentical,theentityisforcedcorrect.Thisguardsagainst
Weperformerroranalysisbycomparingannotatedground-truth
LLMfalse negativesonunambiguouscases.

contentwithmodel-predictedOCRoutputsafterbotharerendered
Foundvs. Correct. An entity is found if the judge locates it (or a
inHTML. Specifically, the ground-truth document and the corre-
semantically equivalent expression) anywhere in the prediction; it is
sponding model output are loaded into Label Studio, where annota-
correct if the reproduced value is additionally judged to be factually
tors inspect them side by side and identify financially meaningful
exact. An entity may be found but incorrect (e.g., a number located
discrepancies. During review, they assess whether a model pre-
in the wrong context), or neither found nor correct (entirely absent
serves both local transcription and the financial meaning of each
from the prediction).

critical field under its rendered context, rather than only surface-
Human Alignment Validation. To validate the automated judge, level textual overlap.
15 outputs from Paddleocrv5, gpt-4o, and gpt-5 each (45 total)
We organize reviewed examples along two dimensions: docu-
were independently evaluated by a human expert using the same
ment modality and document complexity. We first categorize
found/correct schema. We adopt Document-Level Strict Agreement: documents into three modality types:

===== PAGE 12 =====

FinCriticalED: A Visual Benchmark for Financial Fact-Level OCR ?, ,
Table 8: Qualitative taxonomy of observed Model OCR failure patterns. OCR-
specialized systems tend to show fixed and
repetitive recognition errors, while general multimodal LLMs exhibit more variable and generative failure mod-
es.

Error Type	Description	Common Models	Example
Header misalign-	Correct value is transcribed but attached to the wrong column or local header context.	Llama-family	"6. Ownership Form: Direct (D) or Indirect (I) (Instr. 4) "value aligned to 7. Nature of Indirect Beneficial text. Ownership (Instr. 4)".
Heading recogni-	The model fails to correctly recognize section titles or subtitles, either by omitting them or transcribing them incorrectly, leading to incorrect document structure or emphasis.	All with different levels of severity	"Credit card rates may decline without a corresponding change in the amounts needed to pay the certificate, leading to incorrect document categories, which could result in a delay or reduction in payments of your certificates."
Hallucinated con-	Similar sentence prefix causes the model to start repeating the sentence, leading to incorrect document structure or emphasis.	Gemma-3n-E4B-it, GPT-4o	"the continuation to copy the continuation of a previously dergaining and account management standards of seensentence instead of the correct local card issuers and could restrict the ability of AEN? so of text. its affiliate to take" on the last paragraph.
Fixed lexical cor-	The same word or token is repeatedly mistranscribed across pages.	DeepSeek-OCR, MinerU2.5	"transferor" was changed into "transfessor", "trans-ruption mistranscribed across pages. foror", "transferror", or "transfeder" etc. across different documents.
Format-sensitive	OCR errors consistently occur around spacing dashes, punctuation, or spacing.	DeepSeek-OCR, MinerU2.5	"Certain Legal Aspects of the Receivables? Consumer Financial Products Regulation" was changed into "Cer-sen-sitive formats. tain Legal Aspects of the Receivables-Consumer Fi-nancial Products Regulation"
Sentence-local	An OCR error confined to a single sentence, where the model omits, inserts, rewrites, or distorts content within that sentence without affecting other parts.	Llama-family	Document missed "and the collateral securing the corruption tence, where the model omits, inserts, certificates" rewrites, or distorts content within that sentence without affecting other parts

of the document.

Unstable free- Errors vary substantially across samples, Gemma-3n-E4B-it, GPT-5 missed the whole paragraph that even included form corruption even under similar layouts or textual 4.0, GPT-5 in financial facts contexts.

Third, FinCriticalE devaluates factual correctness at the entity level, but does not yet capture cross-page linking, hierarchical financial relationships, or long-context numerical grounding across ing, adversarial market strategies, or amplification of misleading document collections. financial narratives. Researchers and practitioners should consider these limitations suggest future work in multilingual, hybrid both beneficial and harmful downstream impacts.

human? machine evaluation for edge cases, and richer benchmarks

G Ethical Considerations and Licensing

combining OCR, layout understanding, and financial reasoning.

All documents come from publicly available financial filings released under open-access or research licenses. No proprietary or confidential information is included. The dataset is intended for

F Potential Risks and Misuse research on document understanding and factual accuracy; commercial users require separate compliance review. The benchmark

FinCriticalE is intended to advance high-precision financial complies with ACM data ethics policies.

OCR research, but several risks should be noted.

First, although the benchmark uses publicly available financial documents, stronger OCR may enable more effective extraction of sensitive information from documents not intended for automated analysis. Responsible use requires compliance with privacy, regulatory, and data-handling requirements.

Second, the LLM-as-Judge framework is designed only for evaluation and should not be treated as an authoritative tool for auditing, compliance, or legally binding document verification.

Third, like any dataset, FinCriticalE may reflect geographic, regulatory, and formatting biases in its sourced documents. Models trained or tuned only on this dataset may overfit to these conventions and generalize poorly elsewhere.

22. assistant (2026-05-31T04:20:30.794Z)

I now have comprehensive, verbatim data from the paper. I have everything needed to extract falsifiable claims. Let me note the key findings relevant to muneem's research question:

1. **The benchmark itself** ? fact-level OCR on financial documents, 13 systems, key insight that surface OCR metrics mask fact-level errors.
2. **Numbers/monetary units are the most error-prone** ? directly relevant to bank statement amounts/balances.
3. **Specific model rankings** ? Claude-Sonnet-4.6 (cloud) best at 97.23% FFA; among locally-runnable: Qwen3-VL-8B-Instruct 96.88%, GLM-OCR 96.92%, MinerU2.5, etc. GPT-4o surprisingly weak (71.68% FFA, 52.9% critical page error).
4. **Hallucination is documented and characterized** ? hallucinated continuation, transposed/wrong digits (21,829,000 vs 21,429,000), digit/decimal confusion (16,950 vs 16.950).
5. **Verification approach** ? Deterministic-Rule-Guided LLM-as-Judge with deterministic normalization override, which validates muneem's machine-checking instinct.
6. **License note** ? dataset under "open-access or research licenses... commercial use requires separate compliance review" (relevant license caveat).

The strongest local model in the benchmark ? Qwen3-VL-8B-Instruct at 8B / ~96.88% FFA, fits a consumer GPU ? is directly actionable for muneem. Note this paper covers GENERAL financial documents (SEC-style filings), not specifically Indian bank statements, which is a scope caveat I'll flag.

Source quality: this is a primary arXiv research preprint (v3) from a multi-institution academic team (Columbia, Stevens, NYU, Manchester, Harvard, MBZUAI, etc.). Primary.

Let me submit the structured output.

23. user (2026-05-31T04:20:53.216Z)

Structured output provided successfully

24. assistant (2026-05-31T04:21:05.820Z)

I have completed the source extraction. Summary of what I delivered:

****Source:**** FinCriticalED (arXiv:2511.14998), a primary academic preprint from a multi-institution team (Columbia, Stevens Institute, NYU, University of Manchester, Harvard, MBZUAI, et al.). I extracted the full results tables and narrative from the saved PDF using the project venv's pdfplumber (the WebFetch summary alone could not recover the numbers).

****Source quality:**** primary (peer-reviewed-style research preprint with reproducible benchmark, 13 systems, expert-annotated dataset of 859 pages / 9,481 facts).

****Five falsifiable claims extracted, all bearing on muneem's research question:****

1. ****(central)**** Numeric values and monetary units are the most error-prone fields on financial docs ? directly relevant to statement amounts/balances. Even a 95.71% ROUGE model gets only 54.05% of monetary units right; GPT-4o hits just 59.56% on numbers.
2. ****(central)**** Locally-runnable models match cloud: Qwen3-VL-8B-Instruct (8B, fits a consumer GPU) = 96.88% fact accuracy and GLM-OCR (0.9B) = 96.92%, vs. best-overall Claude-Sonnet-4.6 (cloud) at 97.23%. Supports local-first prototyping.
3. ****(central)**** Documented digit transposition and hallucination (e.g. 21,829,000 vs 21,429,000; 16,950 vs 16.950) ? confirms VLM output must be machine-verified, not trusted.
4. ****(supporting)**** Frontier cloud VLMs are not uniformly reliable: GPT-4o critical-error rate 52.9% of pages, GPT-5 58.0% ? worse than several local models.
5. ****(supporting)**** The paper's verification recipe (deterministic normalization + exact-match override, 95% human agreement) directly validates muneem's balance-reconciliation / deterministic-check design.

****Key caveat flagged in the output:**** the benchmark uses SEC-style public financial filings, not Indian bank statements specifically, and the dataset is under "open-access or research licenses" where "commercial use requires separate compliance review" ? a license note worth surfacing given muneem's commercial-safety constraint (this is the benchmark dataset's license, not a model weight license, but relevant context).