

claude-windows / a9fd4426-f234-4e4a-bfb4-1f2c8255c661

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-muneem\a9fd4426-f234-4e4a-bfb4-1f2c8255c661\subagents\workflows\wf_076beba8-59c\agent-a18924e6aacab81b5.jsonl`  
- SHA-256: `7d2491811d198ffe4c831f29df609f98dded6178969f732d92217f110d17ae21`  
- Source modified: `2026-05-31T04:23:52+00:00`  
- Imported at: `2026-07-05T16:48:20+00:00`  
- project: `wf_076beba8-59c`  
- session_id: `a9fd4426-f234-4e4a-bfb4-1f2c8255c661`
```

Transcript

1. user (2026-05-31T04:22:14.996Z)

Adversarial Claim Verifier (voter 2/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Identify the best LOCAL (self-hosted, single NVIDIA CUDA GPU) vision / document-AI / OCR models ? and, as a clearly-flagged approval-gated fallback only, paid cloud API services ? for parsing BANK STATEMENT PDFs into structured transaction tables (date, narration/description, withdrawal/debit, deposit/credit, running balance), for a LOCAL-FIRST personal-finance tool ("muneem").

Target documents: Indian bank statements (HDFC, Axis, ICICI, AU Small Finance) ? dense, multi-column, often SEMI-RULED or borderless tables, sometimes multiple accounts per statement; both born-digital (text layer present) and the harder cases where pdfplumber's extract_tables/coordinate clustering is unreliable (e.g. Axis, whose column layout varies between statements).

HARD CONSTRAINTS (a model that violates these is unusable for us):

1. MUST run 100% LOCALLY for statement contents (highly sensitive). NO cloud APIs by default. Paid/cloud is acceptable ONLY as an explicitly-approved fallback ? still report the best cloud options but label them clearly as "cloud/sensitive ? approval-gated, not default".
2. MUST be usable from Python (transformers / vLLM / Ollama / ONNX Runtime / native lib).
3. License MUST be permissive & commercially safe for BOTH the CODE and the MODEL WEIGHTS: MIT, BSD, Apache-2.0, ISC, HPND. AVOID and explicitly FLAG anything GPL/LGPL/AGPL, CC-BY-NC / non-commercial, research-only, or custom "community"/acceptable-use licenses with restrictions. This is critical ? many document-AI models and their training datasets are non-commercial (e.g. LayoutLMv3 weights are CC-BY-NC; Surya and Marker use GPL-or-commercial dual licensing; several VLM weights ship custom restrictive licenses). Verify the WEIGHTS license, not just the repo license.
4. Hardware: a single consumer/prosumer NVIDIA GPU. Give VRAM tiers (~8-12 GB consumer, ~16-24 GB) and which models/quantizations fit each.

Evaluate and compare these LOCAL approaches ? for each give: code license, MODEL-WEIGHTS license (state explicitly), latest release / maintenance status as of 2025-2026, VRAM need + quantization options, Python integration path, and reliability specifically on DENSE FINANCIAL TABLES / multi-column statements:

- Open vision-language models (VLMs) for document/table understanding: Qwen2-VL and Qwen2.5-VL (note per-size license ? which sizes are Apache-2.0 vs restricted), MiniCPM-V 2.6, InternVL2 / 2.5, Microsoft Phi-3.5-vision, Mistral Pixtral, Meta Llama-3.2-Vision (flag the Llama Community License restrictions), Moondream2, GOT-OCR2.0, olmOCR and other OCR-specialized VLMs.
- OCR-specialist / document-AI pipelines: PaddleOCR + PP-Structure (table recognition), docTR, EasyOCR, Tesseract (already used), Microsoft Table Transformer / TATR (check code AND PubTables-1M dataset/weights license), Donut, Surya (FLAG license), Marker (FLAG license), MinerU / PDF-Extract-Kit, RapidOCR.

- Table-structure recognition specifically (cell/grid reconstruction for borderless tables): TATR, PP-Structure, UniTable, etc.

Also assess: for SEMI-RULED / borderless bank tables, is a prompt-driven VLM (page image -> JSON transactions) more reliable than a classic OCR + table-structure pipeline? What are the accuracy/hallucination risks of VLMs on financial NUMBERS (transposed digits, dropped/duplicated rows), and how should output be VERIFIED? (muneem already machine-checks parses by balance reconciliation ? opening + sum(deposits) - sum(withdrawals) == closing, plus per-row running-balance walks and cross-checks against the bank's printed summary ? so any model's output can be automatically validated, which should inform the recommendation.)

Finally, briefly cover the best PAID/CLOUD fallbacks (clearly flagged approval-gated for sensitive data): AWS Textract (AnalyzeExpense / Tables), Google Document AI, Azure AI Document Intelligence, Reducto, LlamaParse, Anthropic/OpenAI vision ? note any self-hosting option and data-handling terms.

Deliver a RANKED shortlist of the most promising LOCAL options for muneem (NVIDIA GPU, Python, commercial-license-safe), each with license (code + weights), VRAM tier, maturity, strengths/weaknesses for dense bank-statement tables, and a concrete recommendation: which local model(s) to prototype for muneem's statement parsing and how they'd combine with the existing deterministic parsers + balance-reconciliation verification (e.g. VLM as a fallback/cross-check when deterministic parsing fails to reconcile). Cite sources and call out any license uncertainty explicitly.

Claim under review

"Surya's model WEIGHTS are NOT permissively licensed: they use a modified AI Pubs OpenRAIL-M license that is free only for research, personal use, and organizations under \$5M funding/revenue, with a commercial license required above that threshold. This fails muneem's requirement that model weights be MIT/BSD/Apache-2.0 for unrestricted commercial use."

****Source:**** <https://github.com/VikParuchuri/surya/blob/master/LICENSE> (primary)

****Supporting quote:**** "The model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue)"

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-05-31T04:22:17.814Z)

I'll verify this claim about Surya's model weights license. Let me start by examining the primary source and searching for corroborating/contradicting evidence.

3. assistant (2026-05-31T04:22:21.021Z)

Now let me fetch the primary source (the Surya LICENSE file) and run parallel searches for contradicting evidence.

4. user (2026-05-31T04:22:27.878Z)

License Summary for Surya Repository

Based on the displayed LICENSE file, the code is governed by the **Apache License 2.0** (Version 2.0, January 2004).

Key Commercial Use Provisions

The Apache 2.0 license explicitly permits commercial use. The relevant section states:

"You may reproduce and distribute copies of the Work or Derivative Works thereof in any medium, with or without modifications, and in Source or Object form"

There are **no revenue thresholds, commercial use restrictions, non-commercial limitations, or OpenRAIL-M/research-only clauses** mentioned in this license file.

What's Not Present

The displayed content contains only the code license. There is **no information about model weights licensing** in this LICENSE file. Model weight licensing terms would typically be found in separate documentation (often in README files or additional license files in the repository).

The Apache 2.0 license is permissive and allows broad commercial use without funding thresholds or revenue-based restrictions.

5. user (2026-05-31T04:22:28.527Z)

Apache License 2.0 Analysis

Based on my review of the provided text, this is the **Apache License Version 2.0** ? a permissive open source software license.

Key Findings:

No mention of the following terms appear in this license:

- Model weights
- Commercial use restrictions
- Funding/revenue thresholds (\$5M or otherwise)
- OpenRAIL-M
- AI Pubs
- Research-only restrictions
- Non-commercial limitations

What This License Actually Covers:

This Apache 2.0 license grants broad permissions. As stated: "each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable copyright license to reproduce, prepare Derivative Works of, publicly display, publicly perform, sublicense, and distribute the Work."

The license explicitly permits commercial use and derivative works, provided you include proper attribution and license notices with distributed copies.

Summary:

The Apache 2.0 license is fundamentally different from restricted AI model licenses. It contains no AI-specific provisions, model weight regulations, or commercial use barriers. It's a general-purpose open source software license focused on copyright and patent grants with attribution requirements.

6. user (2026-05-31T04:22:31.853Z)

Web search results for query: "Surya OCR VikParuchuri license model weights commercial AI Pubs OpenRAIL-M \$5M"

Links: [{"title": "GitHub - datalab-to/surya: OCR, layout analysis, reading order, table recognition in 90+ languages · GitHub", "url": "https://github.com/datalab-to/surya"}, {"title": "GitHub - CibiQuadance/surya-OCR: OCR and line detection in 90+ languages · GitHub", "url": "https://github.com/CibiQuadance/surya-OCR"}, {"title": "GitHub - cognitiveailab/surya-sci: OCR, layout analysis, reading order, line detection in 90+ languages · GitHub", "url": "https://github.com/cognitiveailab/surya-sci"}, {"title": "surya-ocr · PyPI", "url": "https://pypi.org/project/surya-ocr/"}, {"title": "pitapo/surya · Hugging Face", "url": "https://huggingface.co/pitapo/surya"}, {"title": "Surya OCR | tasarim.ai", "url": "https://tasarim.ai/en/models/surya-ocr"}, {"title": "GitHub - kunthet/surya-text-line-detection-ocr: Accurate line-level text detection and recognition (OCR) in any language · GitHub", "url": "https://github.com/kunthet/surya-text-line-detection-ocr"}, {"title": "cudanexus/ocr-surya | Run with an API on Replicate", "url": "https://replicate.com/cudanexus/ocr-surya"}, {"title": "View as Markdown", "url": "https://replicate.com/cudanexus/ocr-surya/llms.txt"}, {"title": "surya-OCR/README.md at master · CibiQuadance/surya-OCR", "url": "https://github.com/CibiQuadance/surya-OCR/blob/master/README.md"}]

Based on the search results, here's what I found about Surya OCR, its licensing, and commercial AI considerations:

Surya OCR Overview

Surya 2 runs layout, OCR, and table recognition through a single VLM, and it was authored by Vikas Paruchuri and the Datalab Team.

Licensing Model

The model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue). The Surya code is licensed under Apache 2.0.

Commercial Usage

The weights are licensed cc-by-nc-sa-4.0, but the license waives that for any organization under \$5M USD in gross revenue in the most recent 12-month period. For broader commercial licensing of the model weights, visit their pricing page.

The licensing structure is designed to allow free usage for research, personal projects, and small startups with under \$5M in funding or revenue, while requiring commercial licensing agreements for larger organizations or those exceeding the revenue threshold.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

7. assistant (2026-05-31T04:22:36.066Z)

Important nuance emerging: the GitHub LICENSE file itself contains only Apache 2.0 (covering the CODE). The weights-license claim comes from the README, not the LICENSE file. The repo has also moved from `VikParuchuri/surya` to `datalab-to/surya`. Let me verify the actual README text where the weights-license statement lives, and check the dates/current status.

8. user (2026-05-31T04:22:41.034Z)

Surya Licensing Information

Code License

"Code License Apache--2.0" according to the repository badge.

Model Weights License

The README states: "The model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue)."

Commercial Licensing

For broader commercial use of model weights beyond the startup exemption, the documentation directs users to "visit our pricing page" for licensing options.

9. user (2026-05-31T04:22:42.083Z)

Licensing Section

Code License: Apache 2.0

Model Weights License: Modified AI Pubs Open Rail-M

Key Terms from the Documentation

The document states: "The Surya code is licensed under Apache 2.0. The model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue)."

Permitted Uses (Model Weights):

- Research purposes
- Personal use
- Startups with under \$5M in funding or revenue

Commercial Licensing:

For broader commercial licensing beyond the \$5M threshold, the documentation directs users to visit the pricing page at datalab.to/pricing.

10. assistant (2026-05-31T04:22:45.136Z)

The README text is confirmed from the canonical source. Let me verify the actual weights-license file on Hugging Face (where the weights live) to confirm the modified license terms directly, and check the LICENSE_WEIGHTS file in the repo if it exists.

11. user (2026-05-31T04:22:47.304Z)

The server returned HTTP 401 Unauthorized.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

12. user (2026-05-31T04:22:57.350Z)

Web search results for query: "AI Pubs" "OpenRAIL-M" license \$5M revenue commercial Datalab Surya Marker"

Links: [{"title": "GitHub - datalab-to/surya: OCR, layout analysis, reading order, table recognition in 90+ languages · GitHub", "url": "https://github.com/datalab-to/surya"}, {"title": "datalab-to/surya-ocr-2 · Hugging Face", "url": "https://huggingface.co/datalab-to/surya-ocr-2"}, {"title": "surya-ocr · PyPI", "url": "https://pypi.org/project/surya-ocr/"}, {"title": "surya/MODEL_LICENSE at master · datalab-to/surya", "url": "https://github.com/datalab-to/surya/blob/master/MODEL_LICENSE"}, {"title": "surya/README.md at master · datalab-to/surya", "url": "https://github.com/datalab-to/surya/blob/master/README.md"}, {"title": "AI Pubs Open RAIL-M License ? Responsible AI Licenses (RAIL)", "url": "https://www.licenses.ai/ai-pubs-open-railm-vz1"}, {"title": "Datalab: Document AI Foundation Models", "url": "https://idp-software.com/vendors/datalab/"}, {"title": "GitHub - datalab-to/marker: Convert PDF to markdown + JSON quickly with high accuracy · GitHub", "url": "https://github.com/datalab-to/marker"}, {"title": "datalab-to/marker | zdoc.app", "url": "https://www.zdoc.app/en/datalab-to/marker"}, {"title": "https://github.com/datalab-to/surya | Ecosyste.ms: Awesome", "url": "https://awesome.ecosyste.ms/projects/github.com/datalab-to/surya"}]

Based on the search results, I found relevant information about these terms:

Surya's model weights use a modified AI Pubs Open Rail-M license that is free for research, personal use, and startups under \$5M funding/revenue. However, there is a discrepancy in the documentation, as Marker's model weights use a modified AI Pubs Open Rail-M license free for research, personal use, and startups under \$2M funding/revenue.

****Key findings:****

1. ****AI Pubs OpenRAIL-M License****: This is an open machine-learning license with use restrictions prohibiting certain uses of the Model.
2. ****Datalab****: The company behind both Surya (an OCR and layout analysis tool) and Marker (a PDF-to-markdown converter).
3. ****Surya Marker****: These are two separate products developed by Datalab:
 - Surya is a single VLM that handles layout analysis, OCR (full-page or per-block), and table recognition in one model.
 - Marker converts PDFs and other document formats to structured output by combining layout detection, OCR, and optional LLM post-processing.
4. ****Licensing Model****: The GPL code license is paired with a restricted model-weights license that's free below \$2M in revenue and paid above that threshold.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

13. assistant (2026-05-31T04:23:00.319Z)

There's a dedicated MODEL_LICENSE file in the repo. Let me fetch it directly to confirm the actual license text governing the weights, and confirm the AI Pubs OpenRAIL-M license is indeed restrictive.

14. user (2026-05-31T04:23:06.157Z)

AI PUBS OPEN RAIL-M LICENSE (MODIFIED)

****License Name:**** OpenRAIL-M License (Modified), Version 0.1, March 2, 2023

Key Terms on Commercial Restrictions

****\$5M Funding/Revenue Threshold:****

"for any purpose if You (your employer, or the entity you are affiliated with) generated more than two million US Dollars (\$5,000,000) in gross revenue in the prior year"

"for any purpose if You (your employer, or the entity you are affiliated with) has raised more than two million US dollars (\$5,000,000) in total equity or debt funding"

****Research & Personal Use Exception:****

Both restrictions include "except where Your Use is limited to personal use or research purposes"

****Separate License Required:****

Commercial use above these thresholds is restricted unless "Commercial and broader use licenses may be available from Licensor"

****Competitive Use Restriction:****

"for any purpose if You (your employer, or the entity you are affiliated with) provides or otherwise makes available any product or service that competes with any product or service offered by or made available by Licensor"

****Note:**** This is not a CC-BY-NC-SA license. OpenRAIL-M uses custom use-based restrictions

detailed in Attachment A rather than standard Creative Commons terms.

15. user (2026-05-31T04:23:07.633Z)

OpenRAIL-M License Analysis

****Is it permissive like MIT/BSD/Apache?****

No. While OpenRAIL-M grants broad rights to use, modify, and distribute the model, it is fundamentally different from permissive licenses because it includes enforceable use restrictions.

****Does it contain use restrictions?****

Yes. The license explicitly states: "You agree not to use the Model or its Derivatives in any of the following ways" followed by restrictions covering legal compliance, harm prevention, discrimination, and transparency requirements.

****Key restriction examples:****

- Harm-based: "For the purpose of exploiting, Harming or attempting to exploit or Harm minors in any way"
- Discrimination: "For any use intended to or which has the effect of discriminating against individuals or groups based on legally protected characteristics"
- Transparency: "To generate or disseminate machine-generated information or content in any medium without expressly and intelligibly disclaiming that it is machine-generated"

****Is it OSI-approved open-source?****

No. OpenRAIL-M is not an OSI-approved license. The use restrictions violate the Open Source Definition, which requires that licenses permit use for any purpose. The license is categorized as a "Responsible AI License" rather than traditional open-source, representing a newer licensing paradigm focused on responsible model deployment.

16. assistant (2026-05-31T04:23:12.870Z)

I now have direct primary-source confirmation. The MODEL_LICENSE file in the Surya repo explicitly contains the \$5M revenue/funding threshold and commercial-use restriction. Note one minor textual quirk: the modified license document still has the boilerplate phrase "two million US Dollars" next to the numeral "\$5,000,000" ? the numeral is what was changed for Surya's modification, and the README/repo state \$5M for Surya. This is a copy-paste artifact in the license template, not a contradiction of the claim's substance.

Let me do one final cross-check on the claim's precise wording ? specifically whether it's truly "research, personal use, AND organizations under \$5M" ? versus any source disputing this.

17. user (2026-05-31T04:23:19.649Z)

Web search results for query: "Surya OCR weights license non-commercial restriction \$5M revenue commercial license required 2025 2026"

Links: [{"title": "GitHub - CibiQuadance/surya-OCR: OCR and line detection in 90+ languages · GitHub", "url": "https://github.com/CibiQuadance/surya-OCR"}, {"title": "surya-ocr · PyPI", "url": "https://pypi.org/project/surya-ocr/"}, {"title": "GitHub - datalab-to/surya: OCR, layout analysis, reading order, table recognition in 90+ languages · GitHub", "url": "https://github.com/datalab-to/surya"}, {"title": "GitHub - cognitiveailab/surya-sci: OCR, layout analysis, reading order, line detection in 90+ languages · GitHub", "url": "https://github.com/cognitiveailab/surya-sci"}, {"title": "datalab-to/surya-ocr-2 · Hugging Face", "url": "https://huggingface.co/datalab-to/surya-ocr-2"}, {"title": "View as Markdown", "url": "https://replicate.com/cudanexus/ocr-surya/llms.txt"}, {"title": "cudanexus/ocr-surya | Run with an API on Replicate", "url": "https://replicate.com/cudanexus/ocr-surya"}, {"title": "GitHub - adithya-s-k/omniparse: Ingest, parse, and optimize any data format ?? from documents to multimedia ?? for enhanced compatibility with GenAI frameworks", "url": "https://github.com/adithya-s-k/omniparse"}, {"title": "Surya: The Revolutionary

```
OCR Toolkit for 90+ Languages", "url": "https://converter.brightcoding.dev/blog/surya-the-revolutionary-ocr-toolkit-for-90-languages"}, {"title": "Surya", "url": "https://ai4culture.eu/resources/tools/57"}]
```

Based on the search results, here is the information about Surya OCR's licensing and commercial restrictions:

Licensing and Commercial Restrictions

The weights are licensed cc-by-nc-sa-4.0, but the author will waive that for any organization under \$5M USD in gross revenue in the most recent 12-month period.

For the current version (Surya 2), the model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue).

Commercial Licensing Requirements

If you want to use the weights commercially over the revenue limit, please contact surya@vikas.sh for dual licensing. Additionally, for broader commercial licensing of the model weights, visit the pricing page.

Key Points

- **Non-commercial or under \$5M revenue**: The license restrictions are waived
- **Over \$5M in revenue**: Organizations need to contact the author or purchase a commercial license through the pricing page
- **Code License**: The Surya code is licensed under Apache 2.0, which is permissive for commercial use

The licensing terms apply to both the original Surya OCR and the current Surya 2 versions as of 2025-2026.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

18. assistant (2026-05-31T04:23:37.170Z)

I have thorough, converging primary-source evidence. Let me assess the claim against the verification checklist.

Verification summary:

- Supported by quote?** Yes, and verified beyond the quote. The supporting quote ("The model weights use a modified AI Pubs Open Rail-M license (free for research, personal use, and startups under \$5M funding/revenue)") is the verbatim README text from the canonical repo (``datalab-to/surya``, formerly ``VikParuchuri/surya``). I confirmed it at three independent layers:
 - The repo's dedicated `MODEL_LICENSE`` file directly contains the ``$5,000,000`` gross-revenue and equity/debt-funding thresholds, plus a "Commercial and broader use licenses may be available from Licensor" clause and a competitive-use restriction.
 - The README's Licensing section states code = Apache 2.0, weights = modified AI Pubs Open Rail-M.
 - The AI Pubs Open RAIL-M license is independently confirmed (licenses.ai) to be a use-restricted "Responsible AI License" that is **not** OSI-approved and **not** permissive (it violates the Open Source Definition because it restricts uses).
- Contradicting evidence?** None found. Multiple independent sources (README, MODEL_LICENSE, PyPI, secondary writeups) all agree. The only minor discrepancies are: (a) Marker (a sibling Datalab product) uses a \$2M threshold rather than \$5M ? not relevant to Surya; (b) the modified license template has a copy-paste artifact where the spelled-out "two million" wasn't updated to match the "\$5,000,000" numeral ? a drafting quirk, not a substantive contradiction; (c) older descriptions reference cc-by-nc-sa-4.0 while current versions use the AI Pubs Open RAIL-M license ? both are non-permissive and both impose the \$5M waiver structure, so the claim's

substance holds across versions.

3. ****Source quality?**** Primary sources (the repo's own LICENSE/MODEL_LICENSE/README, the license author's own text). Sufficient for the claim's strength.

4. ****Outdated?**** No. Confirmed current for Surya 2 (2025-2026).

5. ****Marketing/speculation?**** No. These are the license terms themselves.

****One precision note on the claim's framing:**** The claim says weights "use a modified AI Pubs OpenRAIL-M license." Strictly, the LICENSE *file* the source URL points to is Apache 2.0 (code only); the weights license is in the separate `MODEL_LICENSE`` file / README. So the cited URL is slightly imprecise about *where* the weights term lives ? but the substantive assertion (weights are non-permissive, \$5M-gated, commercial license required above the threshold, fails MIT/BSD/Apache requirement for unrestricted commercial use) is fully correct and well-evidenced. The claim's conclusion is sound; muneem's CLAUDE.md/memory already flags Surya as license-restricted, consistent with this finding.

The claim survives the refutation attempt.

19. user (2026-05-31T04:23:47.279Z)

Structured output provided successfully

20. assistant (2026-05-31T04:23:51.911Z)

Claim verified as NOT refuted. The assertion that Surya's model weights are non-permissively licensed (modified AI Pubs Open RAIL-M, \$5M revenue/funding-gated, commercial license required above threshold, failing muneem's MIT/BSD/Apache weights requirement) is well-supported by primary sources at three independent layers: the repo README, the dedicated MODEL_LICENSE file (which directly contains the \$5,000,000 thresholds), and the licenses.ai definition confirming AI Pubs Open RAIL-M is a use-restricted, non-OSI license. The only imprecision is that the specific URL cited points to the Apache-2.0 code LICENSE rather than the weights license ? but the weights restriction is confirmed in the same repo, so the claim's substance is sound and current for 2025-2026.