

claude-windows / 42211f80-2169-4025-ad9d-99c1aa29ff70

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\subagents\workflows\wf_6237b110-d58\agent-a7cb7abc3dc8831d6.jsonl`  
- SHA-256: `52809a232fa5d54d72814ea3aded558c828e2cde3a0220bb2c3d0c3a20bbf6c5`  
- Source modified: `2026-06-06T17:36:35+00:00`  
- Imported at: `2026-07-05T16:48:28+00:00`  
- project: `wf_6237b110-d58`  
- session_id: `42211f80-2169-4025-ad9d-99c1aa29ff70`
```

Transcript

1. user (2026-06-06T17:35:16.629Z)

Adversarial Claim Verifier (voter 3/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Research alternative methods to automatically detect rally start/end segments (temporal segmentation) in racket-sport video ? especially badminton ? for a SELF-HOSTED, LOCAL-FIRST pipeline, and assess whether a cloud-VLM teacher ? local-model distillation is the best approach or whether other methods are competitive or complementary. Keep it focused ("mini" research): prioritize actionable, commercially-licensable, locally-runnable approaches over exotic SOTA needing huge compute or restrictive licenses.

OUR CONTEXT / CONSTRAINTS (ground every recommendation in these):

- Self-hosted on a single consumer GPU (RTX 2070 Super), Windows + WSL2. We want LOCAL-ONLY inference at runtime ? privacy lever: user video never leaves the device. It's a commercial product, so PERMISSIVE / commercial-friendly licenses are REQUIRED (avoid AGPL and non-commercial-only weights/datasets).
- Current substrate: WASB (MIT) shuttle-trajectory detector ? per-frame (x, y, visibility). We derive rally windows via velocity/merge-gap/min-duration thresholds plus a net-crossing gate. We ALSO have a permissive person detector available (torchvision Faster R-CNN + supervision ByteTrack) and the video's AUDIO track.
- Current best approach: use paid Gemini (cloud VLM, full-context) OFFLINE as a TEACHER to label rallies, then DISTILL into a local model (logistic regression on shuttle-trajectory features). Findings so far: threshold-tuning ceilings around F1 ? 0.44 and overfits with few videos; the learned model is data-starved at n=2 labeled videos; WASB candidate RECALL is a real bottleneck (its windows sometimes don't align with true rallies at IoU 0.5).
- Inputs we can exploit locally: video frames, the shuttle trajectory, player bounding boxes / pose (if we add a permissive pose model), and the AUDIO track (shuttle "thwack" hits, crowd noise, umpire/score calls).

SPECIFICALLY COVER:

1. Method families for rally/point segmentation in racket sports (badminton, tennis, table tennis, pickleball): audio-based shot/hit detection; temporal action localization (TAL) deep models; pose/player-stance + court-homography state machines (e.g., "serve-ready stance" gating); classical change-point detection / HMM on the shuttle trajectory; scoreboard / score-overlay OCR for point boundaries; self-supervised / weakly-supervised approaches; and VLM teacher?student distillation.
2. For EACH family: how it works, the labels/data it needs, expected accuracy, local-inference COMPUTE cost, and LICENSE / commercial viability of any models, datasets, or weights involved.
3. Badminton rally-segmentation datasets (e.g., ShuttleSet, TrackNet/TrackNetV2 data) and their licenses / commercial usability.
4. A RANKED recommendation for our local-first, permissive, single-GPU setup: is Gemini-teacher ? local-distillation the best, or are some alternatives competitive or complementary ?

especially AUDIO-based hit detection and pose/court-geometry state machines, which we can run fully locally?

5. Concrete, cheap next experiments we could run with what we already have (shuttle trajectory + person detector + audio).

Claim under review

"Precise Event Spotting (PES), the regime needed for frame-accurate boundaries (e.g., shuttle contact / rally edges) in racket sports, requires detection thresholds of roughly 1-5 frames (~33-200 ms), far tighter than the 5-60 second/frame tolerance of general Action Spotting."

****Source:**** <https://arxiv.org/html/2505.03991v1> (primary)
****Supporting quote:**** "PES requires extremely strict detection thresholds, typically constrained within 1?5 frames (approximately 33?200 milliseconds)"

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.
****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.
Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-06-06T17:35:20.036Z)

I'll verify this claim about Precise Event Spotting (PES) detection thresholds. Let me start by loading the necessary tools and examining the primary source.

3. assistant (2026-06-06T17:35:22.873Z)

Now let me fetch the primary source to verify the exact quote and context.

4. user (2026-06-06T17:35:32.775Z)

Survey Paper Analysis

****Title:**** Action Spotting and Precise Event Detection in Sports: Datasets, Methods, and Challenges

****Authors:**** Hao Xu, Arbind Agrahari Baniya, Sam Wells, Mohamed Reda Bouadjenek, Richard Dazeley, Sunil Aryal

****Date:**** May 6, 2025

****Subject:**** A comprehensive review of deep learning approaches for video event detection in sports, focusing on Temporal Action Localization (TAL), Action Spotting (AS), and Precise Event Spotting (PES).

PES Detection Thresholds and Temporal Precision

****Exact Quote on Frame Tolerance:****
"PES requires extremely strict detection thresholds, typically constrained within 1?5 frames

(approximately 33?200 milliseconds at standard frame rates of 25?30 FPS)."

****Context:**** This statement appears in Section 2.3 (Precise Event Spotting) and establishes PES's core distinction from Action Spotting.

****Comparison to Action Spotting:****

The document provides direct comparison in Table 1, showing:

- Action Spotting tolerance: "5?60 frames"
- PES tolerance: "0?2 frames"

****Additional Context Statement:****

"While AS facilitates coarser detection by tolerating errors within relatively broad temporal windows (typically between 5 and 60 seconds), certain applications in sports demand much higher temporal precision."

This explicitly contrasts the two approaches, establishing that PES operates at substantially tighter tolerances than general Action Spotting methods.

5. assistant (2026-06-06T17:35:36.321Z)

The primary source supports the claim well. Let me verify a couple of details ? the exact framing of the Action Spotting tolerance (the claim says "5-60 second/frame tolerance") and cross-check against independent sources on PES thresholds to ensure this isn't an isolated or idiosyncratic claim.

6. user (2026-06-06T17:35:43.757Z)

Action Spotting vs Precise Event Spotting: Tolerance Windows

****Action Spotting (AS) Tolerance:****

The document states: "AS facilitates coarser detection by tolerating errors within relatively broad temporal windows (typically between 5 and 60 seconds)"

Additionally, Table 1 specifies AS uses a "5?60 frames" tolerance window.

****Precise Event Spotting (PES) Tolerance:****

The document indicates: "PES was proposed as an extension of AS. PES requires extremely strict detection thresholds, typically constrained within 1?5 frames (approximately 33?200 milliseconds at standard frame rates of 25?30 FPS)"

Table 1 confirms PES uses a "0?2 frames" tolerance window.

****Hong et al. on PES Introduction:****

The paper references: "J. Hong et al. later extended AS to PES by setting a stricter tolerance window" and notes they "proposed E2E-Spot model" with frame-level precision requirements, though the exact threshold values from their original work are presented in the survey's summary rather than as direct quotations from Hong et al.'s paper.

7. user (2026-06-06T17:35:47.988Z)

Web search results for query: "Precise Event Spotting PES frame tolerance 1-5 frames temporal precision sports video"

Links: [{"title": "Multi-Focus Temporal Shifting for Precise Event Spotting in Sports Videos", "url": "https://arxiv.org/html/2507.07381"}, {"title": "Few-Shot Precise Event Spotting via Unified Multi-Entity Graph and Distillation", "url": "https://arxiv.org/html/2511.14186"}, {"title": "Precise Event Spotting in Sports Videos: Solving Long-Range Dependency and Class Imbalance", "url": "https://arxiv.org/html/2503.00147v1"}, {"title": "Few-Shot Precise Event Spotting via Unified Multi-Entity Graph and Distillation", "url": "https://arxiv.org/pdf/2511.14186"}, {"title": "[2507.07381] Multi-Focus Temporal Shifting for Precise Event Spotting in Sports Videos", "url": "https://arxiv.org/abs/2507.07381"}, {"title": "Precise Event Spotting in Sports Videos: Solving Long-Range Dependency and Class Imbalance | Request PDF", "url": "https://www.rese"}]

archgate.net/publication/394629549_Precise_Event_Spotting_in_Sports_Videos_Solving_Long-Range_Dependency_and_Class_Imbalance"}, {"title": "AdaSpot: Spend Resolution Where It Matters for Precise Event Spotting", "url": "https://arxiv.org/html/2602.22073"}, {"title": "Precise Event Spotting in Sports Videos: Solving Long-Range Dependency and Class Imbalance", "url": "https://arxiv.org/pdf/2503.00147"}, {"title": "Action Spotting and Precise Event Detection in Sports: Datasets, Methods, and Challenges", "url": "https://arxiv.org/html/2505.03991v1"}, {"title": "Spotting Temporally Precise, Fine-Grained Events in Video", "url": "https://arxiv.org/pdf/2207.10213"}]

Based on the search results, here's what I found about Precise Event Spotting (PES), frame tolerance, and temporal precision in sports videos:

Definition and Frame Tolerance

Precise Event Spotting (PES) aims to localize and classify events at frame-level accuracy, typically within a 1?2 frame tolerance. More specifically, PES imposes a much stricter tolerance, typically constrained to 0?2 frames, requiring frame-level temporal precision.

Why This Precision Matters in Sports

This fine-grained precision is critical in sports scenarios, where even a few frames of deviation can correspond to a different event entirely. For example, in sports such as tennis and figure skating, even a margin of error of 1?2 frames can result in missing critical events, such as the precise moment of ball contact in tennis.

Applications and Benefits

Accurate PES provides a strong foundation for downstream tasks such as ball tracking, player positioning, and automatic highlight generation.

Evaluation Methodology

Both action spotting and PES treat detections as successful if the predicted timestamp falls within a predefined tolerance window centered around the ground-truth event. In action spotting, this window typically ranges from 5 to 60 seconds, while PES imposes a much stricter tolerance, typically constrained to 0?2 frames.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

8. assistant (2026-06-06T17:35:51.366Z)

The claim is corroborated by multiple independent sources. Let me verify the foundational PES paper (Hong et al., "Spotting Temporally Precise, Fine-Grained Events in Video") to confirm the original threshold definition and check for any qualification or contradiction.

9. user (2026-06-06T17:35:57.536Z)

I cannot extract the specific information about frame tolerance thresholds from this PDF document. While the content appears to be from a computer vision/video analysis paper (based on the JPEG image data and technical references visible in the binary PDF structure), the text is heavily compressed and encoded in the PDF stream, making it impossible for me to reliably quote the exact tolerance values or methodology used for evaluation.

To answer your question about precise event spotting tolerance thresholds and how they contrast with SoccerNet's second-based tolerance, I would need either:

- A text-readable version of the paper
- The uncompressed content clearly visible in the source

I recommend accessing the paper directly through its official source to find sections discussing evaluation metrics and temporal tolerance specifications.

[Binary content (application/pdf, 9.8MB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\tool-results\webfetch-1780767317460-0vhbu5.pdf]

10. assistant (2026-06-06T17:36:00.908Z)

The PDF binary couldn't be parsed. I already have strong corroboration from the survey paper plus a second independent paper (arxiv 2503.00147) confirming the same numbers. Let me do one final cross-check on the foundational paper's tolerance via search to confirm the original Hong et al. threshold.

11. user (2026-06-06T17:36:09.062Z)

Web search results for query: "Hong "Spotting Temporally Precise Fine-Grained Events" E2E-Spot tolerance frames evaluation metric tennis figure skating"

Links: [{"title": "Few-Shot Precise Event Spotting via Unified Multi-Entity Graph and Distillation", "url": "https://arxiv.org/html/2511.14186"}, {"title": "GitHub - jhong93/spot: Code for Spotting Temporally Precise, Fine-Grained Events in Video", "url": "https://github.com/jhong93/spot"}, {"title": "Few-Shot Precise Event Spotting via Unified Multi-Entity Graph and Distillation", "url": "https://arxiv.org/pdf/2511.14186"}, {"title": "(PDF) Few-Shot Precise Event Spotting via Unified Multi-Entity Graph and Distillation", "url": "https://www.researchgate.net/publication/397642798_Few-Shot_Precise_Event_Spotting_via_Unified_Multi-Entity_Graph_and_Distillation"}, {"title": "Multi-Focus Temporal Shifting for Precise Event Spotting in Sports Videos", "url": "https://arxiv.org/html/2507.07381"}, {"title": "Spotting Temporally Precise, Fine-Grained Events in Video", "url": "https://arxiv.org/pdf/2207.10213"}, {"title": "Deep learning for action spotting in association football videos", "url": "https://arxiv.org/pdf/2410.01304"}, {"title": "[2207.10213] Spotting Temporally Precise, Fine-Grained Events in Video", "url": "https://arxiv.org/abs/2207.10213"}, {"title": "Spotting Temporally Precise, Fine-Grained Events in Video | Springer Nature Link", "url": "https://link.springer.com/chapter/10.1007/978-3-031-19833-5_3"}, {"title": "Spotting Temporally Precise, Fine-Grained Events in Video | Request PDF", "url": "https://www.researchgate.net/publication/365099624_Spotting_Temporally_Precise_Fine-Grained_Events_in_Video"}]

Based on the search results, here's what I found about E2E-Spot and the related research:

E2E-Spot Framework

E2E-Spot is an end-to-end deep learning framework for temporally precise spotting of fine-grained events in video, defined as predicting the exact frame (within 1?2 frames tolerance) when an event occurs.

Task Definition and Evaluation Metric

A prediction is deemed correct if its timestamp falls within ? frames of a labeled ground-truth event and it has the correct class label. In precise spotting, the temporal tolerance ? is small ? i.e., a few frames only. Precise Event Spotting (PES) aims to localize and classify events at frame-level accuracy, typically within a 1?2 frame tolerance.

Sports Applications

The framework has been successfully applied to fine-grained sports event detection. E2E-Spot (hong2022spotting), employing a CNN backbone with Gated Shift Modules and GRUs for temporal modeling, performed effectively in tennis and figure skating scenarios.

Model Architecture

Unlike prior two-stage methods (feature extraction followed by separate head training), E2E-Spot jointly learns spatial?temporal representations directly from raw pixels under end-to-end

supervision, enabling both fine-grained motion sensitivity and global temporal reasoning.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

12. assistant (2026-06-06T17:36:21.984Z)

I have sufficient evidence to render a verdict. Let me assess the claim against the checklist.

The claim states: "PES requires detection thresholds of roughly 1-5 frames (~33-200 ms), far tighter than the 5-60 second/frame tolerance of general Action Spotting."

****Findings:****

1. ****Quote support****: The primary source (arXiv 2505.03991v1, a peer-style survey, May 2025) states verbatim: "PES requires extremely strict detection thresholds, typically constrained within 1?5 frames (approximately 33?200 milliseconds at standard frame rates of 25?30 FPS)." This is an exact match to the claim's numbers ? not an overreach.
2. ****Action Spotting contrast****: The source confirms AS tolerates "errors within relatively broad temporal windows (typically between 5 and 60 seconds)." The claim's slightly awkward "5-60 second/frame" phrasing reflects an ambiguity in the source itself ? Table 1 lists AS as "5-60 frames" while the prose says "5-60 seconds." The dominant/standard definition (SoccerNet AS) is seconds, and the claim captures the directionally correct point that PES is "far tighter." The frames-vs-seconds wording is a minor imprecision in the source, not a fabrication by the claim.
3. ****Independent corroboration****: A second paper (arXiv 2503.00147) independently states AS "typically ranges from 5 to 60 seconds, while PES imposes a much stricter tolerance, typically constrained to 0?2 frames." The foundational Hong et al. paper (E2E-Spot, arXiv 2207.10213) defines correctness as falling within ? frames where ? is "a few frames only" (1-2 frames). Multiple independent sources agree.
4. ****Currency****: Sources are 2022-2025 ? current for this fast-moving field.
5. ****Source quality****: Academic survey + foundational paper + multiple corroborating arXiv papers. Not marketing/forum/press-release.

The only quibble is the "5-60 second/frame" wording, which is internally messy ? but it traces directly to an inconsistency in the cited primary source (prose says seconds, table says frames), and the core quantitative claim (PES = 1-5 frames / 33-200ms, far tighter than AS) is well-supported and independently corroborated. This does not meet the bar for refutation.

13. user (2026-06-06T17:36:31.174Z)

Structured output provided successfully

14. assistant (2026-06-06T17:36:35.848Z)

Verdict submitted: ****refuted=false**** (confidence: high).

The claim is well-supported by its primary source and independently corroborated by multiple academic papers (2022-2025). The exact numbers (PES = 1-5 frames / ~33-200ms, far tighter than Action Spotting's seconds-scale tolerance) match verbatim quotes. The only blemish is the claim's "5-60 second/frame" phrasing, which inherits an internal inconsistency from the cited source's own table-vs-prose discrepancy ? not enough to refute a multiply-corroborated, current, academically-sourced claim.