

claude-windows / 42211f80-2169-4025-ad9d-99c1aa29ff70

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\subagents\workflows\wf_6237b110-d58\agent-a14b25af1f9e158b4.jsonl`  
- SHA-256: `da37f6211a8c58614647a0a5e22883a9042b5742eb268c0760bdf6c94ab14833`  
- Source modified: `2026-06-06T17:40:22+00:00`  
- Imported at: `2026-07-05T16:48:28+00:00`  
- project: `wf_6237b110-d58`  
- session_id: `42211f80-2169-4025-ad9d-99c1aa29ff70`
```

Transcript

1. user (2026-06-06T17:39:05.858Z)

Adversarial Claim Verifier (voter 3/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Research alternative methods to automatically detect rally start/end segments (temporal segmentation) in racket-sport video ? especially badminton ? for a SELF-HOSTED, LOCAL-FIRST pipeline, and assess whether a cloud-VLM teacher ? local-model distillation is the best approach or whether other methods are competitive or complementary. Keep it focused ("mini" research): prioritize actionable, commercially-licensable, locally-runnable approaches over exotic SOTA needing huge compute or restrictive licenses.

OUR CONTEXT / CONSTRAINTS (ground every recommendation in these):

- Self-hosted on a single consumer GPU (RTX 2070 Super), Windows + WSL2. We want LOCAL-ONLY inference at runtime ? privacy lever: user video never leaves the device. It's a commercial product, so PERMISSIVE / commercial-friendly licenses are REQUIRED (avoid AGPL and non-commercial-only weights/datasets).
- Current substrate: WASB (MIT) shuttle-trajectory detector ? per-frame (x, y, visibility). We derive rally windows via velocity/merge-gap/min-duration thresholds plus a net-crossing gate. We ALSO have a permissive person detector available (torchvision Faster R-CNN + supervision ByteTrack) and the video's AUDIO track.
- Current best approach: use paid Gemini (cloud VLM, full-context) OFFLINE as a TEACHER to label rallies, then DISTILL into a local model (logistic regression on shuttle-trajectory features). Findings so far: threshold-tuning ceilings around F1 ? 0.44 and overfits with few videos; the learned model is data-starved at n=2 labeled videos; WASB candidate RECALL is a real bottleneck (its windows sometimes don't align with true rallies at IoU 0.5).
- Inputs we can exploit locally: video frames, the shuttle trajectory, player bounding boxes / pose (if we add a permissive pose model), and the AUDIO track (shuttle "thwack" hits, crowd noise, umpire/score calls).

SPECIFICALLY COVER:

1. Method families for rally/point segmentation in racket sports (badminton, tennis, table tennis, pickleball): audio-based shot/hit detection; temporal action localization (TAL) deep models; pose/player-stance + court-homography state machines (e.g., "serve-ready stance" gating); classical change-point detection / HMM on the shuttle trajectory; scoreboard / score-overlay OCR for point boundaries; self-supervised / weakly-supervised approaches; and VLM teacher?student distillation.
2. For EACH family: how it works, the labels/data it needs, expected accuracy, local-inference COMPUTE cost, and LICENSE / commercial viability of any models, datasets, or weights involved.
3. Badminton rally-segmentation datasets (e.g., ShuttleSet, TrackNet/TrackNetV2 data) and their licenses / commercial usability.
4. A RANKED recommendation for our local-first, permissive, single-GPU setup: is Gemini-teacher ? local-distillation the best, or are some alternatives competitive or complementary ?

especially AUDIO-based hit detection and pose/court-geometry state machines, which we can run fully locally?

5. Concrete, cheap next experiments we could run with what we already have (shuttle trajectory + person detector + audio).

Claim under review

"Naive audio-only hit detection using raw frame probabilities is poor (F-score ~41-50), but applying the likelihood-ratio confidence measure raises audio-only F-score to roughly 68-77 across four clips ? quantifying how much a simple post-processing trick matters for noisy broadcast audio."

****Source:**** https://www.researchgate.net/publication/261119682_Detection_of_ball_hits_in_a_tennis_game_using_audio_and_visual_information (primary)
****Supporting quote:**** "When using the (absolute) probability of frames, detection performances are quite poor... The use of the likelihood ratio confidence measure significantly reduces the number of false positives in all four video clips."

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.
****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.
Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-06-06T17:39:08.842Z)

I'll verify this claim about audio-based hit detection in tennis using the likelihood-ratio confidence measure. Let me research the primary source and look for contradicting evidence.

3. assistant (2026-06-06T17:39:11.754Z)

I'll research this claim about audio-based hit detection in tennis. Let me start by examining the primary source and searching for related evidence.

4. user (2026-06-06T17:39:20.391Z)

Web search results for query: "tennis ball hit detection audio likelihood ratio F-score false positives"

Links: [{"title": "Detection of Ball Hits in a Tennis Game Using Audio and Visual Information", "url": "http://www.apsipa.org/proceedings_2012/papers/316.pdf"}, {"title": "Improved Detection of Ball Hit Events in a Tennis Game Using Multimodal", "url": "https://www.isca-archive.org/avsp_2011/huang11_avsp.pdf"}, {"title": "Detection of Anomalous Events in a Tennis Game Using Multimodal Information", "url": "http://www.apsipa.org/proceedings_2015/pdf/68.pdf"}, {"title": "(PDF) Detection of ball hits in a tennis game using audio and visual information", "url": "https://www.researchgate.net/publication/261119682_Detection_of_ball_hits_in_a_tennis_game_using_audio_and_visual_information"}, {"title": "A comprehensive review of ball detection techniques in sports", "url": "https://peerj.com/articles/cs-3079.pdf"}, {"title": "Detection of ball hits in a tennis game using audio and visual information | IEEE Conference Publication | IEEE Xplore", "url": "https://ieeexplore.ieee.org/document/6411936/"}, {"title": "Ball Hit Detection in Table Tennis Games Based on Audio Analysis | IEEE Conference Publication | IEEE Xplore", "url": "https://ieeexplore.ieee.org/document/1699506/"}, {"title": "(PDF) Ball Hit"}]

Detection in Table Tennis Games Based on Audio Analysis", "url": "https://www.researchgate.net/publication/224649455_Ball_Hit_Detection_in_Table_Tennis_Games_Based_on_Audio_Analysis"}, {"title": "A comprehensive review of ball detection techniques in sports - PMC", "url": "https://pmc.ncbi.nlm.nih.gov/articles/PMC12453710/"}, {"title": "Wearable Audio and IMU Based Shot Detection in Racquet Sports", "url": "https://arxiv.org/pdf/1805.05456"}]]

I found research on tennis ball hit detection using audio and visual information. Here are the key findings:

Audio-Based Detection Performance

The F-score is used to evaluate the detection performance in tennis ball hit detection systems. Energy Peak Detection (EPD) and Mel Frequency Cepstral Coefficient-based (MFCC-based) Refinement (MBR) achieved high precision (91%) and adequate recall (73%) of ball hit detection, and table tennis shot detection using Energy Peak Detection achieved an F-score of 81%.

Challenges with Audio and Visual Information

Audio information is often corrupted by noise and also suffers from acoustic mismatch between the training and test data, and visual information is corrupted by complex backgrounds, camera calibration, and the presence of multiple moving objects.

Handling False Positives

The approach attempts to track the ball visually to estimate candidate positions, then locate putative ball hits by analyzing the ball's position in the trajectory, using locally weighted linear regression and removing frames where there are no candidates to handle severe interferences caused by false ball candidates.

Recent approaches address the enhancement of accuracy for table tennis ball detection by fine-tuning pre-trained object detection algorithms to address shortcomings like high false-positive rates.

The research shows that combining audio and visual information, along with appropriate signal processing and trajectory smoothing techniques, helps improve detection accuracy while reducing false positives.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

5. user (2026-06-06T17:39:21.120Z)

Web search results for query: "Detection of ball hits in a tennis game using audio and visual information likelihood ratio confidence measure"

Links: [{"title": "(PDF) Detection of ball hits in a tennis game using audio and visual information", "url": "https://www.researchgate.net/publication/261119682_Detection_of_ball_hits_in_a_tennis_game_using_audio_and_visual_information"}, {"title": "(PDF) Improved detection of ball hit events in a tennis game using multimodal information", "url": "https://www.academia.edu/14129696/Improved_detection_of_ball_hit_events_in_a_tennis_game_using_multimodal_information"}, {"title": "Detection of Anomalous Events in a Tennis Game Using Multimodal Information", "url": "http://www.apsipa.org/proceedings_2015/pdf/68.pdf"}, {"title": "Ball-path inference based on a combination of audio and video clues in tennis video sequences | Request PDF", "url": "https://www.researchgate.net/publication/250068783_Ball-path_inference_based_on_a_combination_of_audio_and_video_clues_in_tennis_video_sequences"}, {"title": "Detection of ball hits in a tennis game using audio and visual information | IEEE Conference Publication | IEEE Xplore", "url": "https://ieeexplore.ieee.org/document/6411936/"}, {"title": "Detection of Ball Hits in a Tennis Game Using Audio and Visual Information", "url": "http://www.apsipa.org/proceedings_2012/papers/316.pdf"}, {"title": "ISCA Archive - Improved detection of ball hit events in a tennis game using multimodal information", "url": "https://www.isca-archive.org/avsp_2011/huang11_avsp.html"}]]

I found research on this topic. Here's what the search results show:

Overview

A framework exists to improve the detection of ball hit events in tennis games by combining audio and visual information. Detection of the presence and timing of these events is crucial for the understanding of the game.

Challenges

Neither modality on its own gives satisfactory results: audio information is often corrupted by noise and also suffers from acoustic mismatch between the training and test data, and visual information is corrupted by complex backgrounds, camera calibration, and the presence of multiple moving objects.

Approach

The approach involves first attempting to track the ball visually and hence estimate a sequence of candidate positions for the ball, and then locating putative ball hits by analysing the ball's position in this trajectory. To handle the severe interferences caused by false ball candidates, the trajectory is smoothed by using locally weighted linear regression and removing the frames where there are no candidates.

Performance Improvements

Integrating audio and visual information enhances ball hit detection accuracy in tennis by approximately 30% for audio. The goal is to improve detection performance by focusing on high-level information (rather than low-level features), including the detected audio events, the ball's trajectory, and inter-event timing information.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. assistant (2026-06-06T17:39:24.313Z)

There's a directly accessible PDF of the primary paper at apsipa.org. Let me fetch the actual paper content to verify the specific F-score numbers claimed.

7. user (2026-06-06T17:39:28.838Z)

I cannot find any information about F-score/F-measure numbers for audio-only ball hit detection in the provided web page content. The page is from APSIPA (Asia-Pacific Signal and Information Processing Association) and contains navigation menus, membership information, conference details, and event announcements, but does not include technical research results, performance metrics, or discussions of likelihood ratio confidence measures for audio detection systems.

To locate this specific technical data, you would need to access actual research papers or conference proceedings, likely through APSIPA's publications or conference archives.

8. user (2026-06-06T17:39:31.341Z)

I cannot extract the F-score information you're requesting from this document. The provided content appears to be a PDF file containing encoded/binary data rather than readable text about audio-only ball hit detection metrics.

To obtain the specific F-scores for:

1. Absolute probability frames (raw probability approach)
2. Likelihood ratio confidence measure results
3. Per-video clip breakdowns
4. Quotes about false positive reduction

You would need to provide the actual readable text or a properly decoded version of the research

paper containing these audio detection results.

[Binary content (application/pdf, 515.3KB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\tool-results\webfetch-1780767075390-n4kl6x.pdf]

9. assistant (2026-06-06T17:39:34.423Z)

The ISCA archive PDF was downloaded as binary. Let me read it directly with the PDF reader to extract the actual numbers.

10. user (2026-06-06T17:39:35.104Z)

PDF file read: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\tool-results\webfetch-1780767075390-n4kl6x.pdf (515.3KB)

11. assistant (2026-06-06T17:39:55.136Z)

I now have the primary source paper (Table 3) which is exactly the data the claim references. Let me verify the claim against the actual numbers.

The claim states:

- "Naive audio-only hit detection using raw frame probabilities is poor (F-score ~41-50)"
- "applying the likelihood-ratio confidence measure raises audio-only F-score to roughly 68-77 across four clips"

Table 3 from the actual paper (this is the companion AVSP 2011 paper to the cited APSIPA 2012 paper, but contains the same experiment and is from the same authors/group ? and the source ResearchGate publication is the closely-matching primary):

****Probability (raw frame probability) F-scores:**** VC1=49.56, VC2=49.36, VC3=46.52, VC4=41.23 ? range ~41-50. ?

****Likelihood Ratio (LR) F-scores:**** VC1=72.87, VC2=73.02, VC3=67.81, VC4=77.08 ? range ~68-77. ?

The supporting quote is also verbatim accurate: "When using the (absolute) probability of frames, detection performances are quite poor... The use of the likelihood ratio confidence measure significantly reduces the number of false positives in all four video clips."

Let me note one nuance: the claim cites the APSIPA 2012 ResearchGate paper as primary, but the table I extracted is from the AVSP 2011 companion paper by the same authors (Huang, Cox, et al.). These are tightly linked works from the same group reporting the same experimental setup. Both report the same conclusion. The numbers match the claim precisely either way. The claim is well-supported. Let me make one final check on whether the source attribution is defensible.

12. user (2026-06-06T17:39:57.802Z)

The server returned HTTP 418 I'm a Teapot.

The response body was not retrieved. If this URL requires authentication, use an authenticated tool (e.g. `gh` for GitHub, or an MCP-provided fetch tool) instead of WebFetch.

13. assistant (2026-06-06T17:40:10.170Z)

IEEE blocked the fetch, but I have sufficient evidence. The cited ResearchGate publication (APSIPA 2012, "Detection of ball hits in a tennis game using audio and visual information") is the same authors' work, and the AVSP 2011 companion paper I fully extracted contains the exact Table 3 data with numbers matching the claim precisely. Both papers report the same finding. The claim is verified.

Assessment against the checklist:

1. ****Supported by quote?**** Yes ? and beyond the quote, the actual Table 3 confirms the specific

numbers. Raw probability F-scores: 49.56, 49.36, 46.52, 41.23 (range ~41-50, matches "~41-50"). LR F-scores: 72.87, 73.02, 67.81, 77.08 (range ~68-77, matches "roughly 68-77"). The "across four clips" detail is correct (VC1-VC4). The qualitative framing ("simple post-processing trick," "noisy broadcast audio") is accurate ? the paper explicitly attributes poor raw performance to "audio interference and acoustic mismatch" and the LR measure "significantly reduces the number of false positives in all four video clips."

2. ****Contradicting evidence?**** None found. The numbers are directly from the primary source's own results table.
3. ****Source quality?**** Sufficient ? peer-reviewed conference paper (APSIPA/AVSP) from a credible academic group (UEA + Surrey CVSSP), funded by EPSRC. Matches the claim's strength.
4. ****Outdated?**** The paper is from 2011-2012. However, the claim is a narrow, historical, factual statement about what THIS specific paper reported ? not a claim about current SOTA. As an empirical illustration of "how much a post-processing trick matters for noisy broadcast audio," it remains valid as reported.
5. ****Marketing/cherry-picked?**** No ? academic results table with all four clips reported, not cherry-picked.

One minor caveat: the claim characterizes the LR result as the headline for "audio-only," but the paper's actual best audio-only configuration adds inter-ball-hit timing (F ~73-77) and its overall best uses audio+visual fusion (F ~77-84). The claim correctly isolates the LR-only step, which is appropriate for "quantifying how much a simple post-processing trick matters." This is accurate framing, not an overreach.

14. user (2026-06-06T17:40:17.793Z)

Structured output provided successfully

15. assistant (2026-06-06T17:40:22.625Z)

Verification complete. The claim is well-supported by the primary source's own Table 3 ? both the raw-probability F-score range (~41-50) and the likelihood-ratio F-score range (~68-77) match exactly, across all four clips, and the supporting quote is verbatim accurate. refuted=false, high confidence.