

claude-windows / 42211f80-2169-4025-ad9d-99c1aa29ff70

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\subagents\workflows\wf_6237b110-d58\agent-aa352e1976d1f6338.jsonl`  
- SHA-256: `7a89da523bea057fafff2cce58afeb404b4f623dc95a12e1f856bea16cce0df1`  
- Source modified: `2026-06-06T17:39:19+00:00`  
- Imported at: `2026-07-05T16:48:28+00:00`  
- project: `wf_6237b110-d58`  
- session_id: `42211f80-2169-4025-ad9d-99c1aa29ff70`
```

Transcript

1. user (2026-06-06T17:37:38.402Z)

Adversarial Claim Verifier (voter 3/3)

Be SKEPTICAL. Try to REFUTE this claim. ?2/3 refutations kill it.

Research question

Research alternative methods to automatically detect rally start/end segments (temporal segmentation) in racket-sport video ? especially badminton ? for a SELF-HOSTED, LOCAL-FIRST pipeline, and assess whether a cloud-VLM teacher ? local-model distillation is the best approach or whether other methods are competitive or complementary. Keep it focused ("mini" research): prioritize actionable, commercially-licensable, locally-runnable approaches over exotic SOTA needing huge compute or restrictive licenses.

OUR CONTEXT / CONSTRAINTS (ground every recommendation in these):

- Self-hosted on a single consumer GPU (RTX 2070 Super), Windows + WSL2. We want LOCAL-ONLY inference at runtime ? privacy lever: user video never leaves the device. It's a commercial product, so PERMISSIVE / commercial-friendly licenses are REQUIRED (avoid AGPL and non-commercial-only weights/datasets).
- Current substrate: WASB (MIT) shuttle-trajectory detector ? per-frame (x, y, visibility). We derive rally windows via velocity/merge-gap/min-duration thresholds plus a net-crossing gate. We ALSO have a permissive person detector available (torchvision Faster R-CNN + supervision ByteTrack) and the video's AUDIO track.
- Current best approach: use paid Gemini (cloud VLM, full-context) OFFLINE as a TEACHER to label rallies, then DISTILL into a local model (logistic regression on shuttle-trajectory features). Findings so far: threshold-tuning ceilings around F1 ? 0.44 and overfits with few videos; the learned model is data-starved at n=2 labeled videos; WASB candidate RECALL is a real bottleneck (its windows sometimes don't align with true rallies at IoU 0.5).
- Inputs we can exploit locally: video frames, the shuttle trajectory, player bounding boxes / pose (if we add a permissive pose model), and the AUDIO track (shuttle "thwack" hits, crowd noise, umpire/score calls).

SPECIFICALLY COVER:

1. Method families for rally/point segmentation in racket sports (badminton, tennis, table tennis, pickleball): audio-based shot/hit detection; temporal action localization (TAL) deep models; pose/player-stance + court-homography state machines (e.g., "serve-ready stance" gating); classical change-point detection / HMM on the shuttle trajectory; scoreboard / score-overlay OCR for point boundaries; self-supervised / weakly-supervised approaches; and VLM teacher?student distillation.
2. For EACH family: how it works, the labels/data it needs, expected accuracy, local-inference COMPUTE cost, and LICENSE / commercial viability of any models, datasets, or weights involved.
3. Badminton rally-segmentation datasets (e.g., ShuttleSet, TrackNet/TrackNetV2 data) and their licenses / commercial usability.
4. A RANKED recommendation for our local-first, permissive, single-GPU setup: is Gemini-teacher ? local-distillation the best, or are some alternatives competitive or complementary ?

especially AUDIO-based hit detection and pose/court-geometry state machines, which we can run fully locally?

5. Concrete, cheap next experiments we could run with what we already have (shuttle trajectory + person detector + audio).

Claim under review

"Audio cues (specifically ball hits plus end-of-exchange applause) are discriminative for rally vs. non-rally scenes, and fusing audio with visual cues in an HMM improves rally classification over either modality alone."

****Source:**** <https://www.irisa.fr/metiss/gravier/biblio/03/kijak-icme03.pdf> (primary)

****Supporting quote:**** "This suggests that audio features are effective to describe rally scenes. Indeed, a rally is essentially characterized by the presence of ball hits and applause which happen at the end of the exchange, although a missed first serve is only characterized by the presence of ball hits... integrating multiple cues from different media improves performance. In particular, ball hits detection improves classification rates for rallies and first serves."

Checklist

1. Is the claim actually supported by the quote, or is it an overreach/misread?
2. WebSearch for contradicting evidence ? does any credible source dispute or heavily qualify this?
3. Is the source quality sufficient for the claim's strength? (extraordinary claims need primary sources)
4. Is the claim outdated? (check dates ? old claims about fast-moving fields are suspect)
5. Is this a marketing claim / press release / cherry-picked benchmark / forum speculation?

****refuted=true**** if: unsupported by quote / contradicted / low-quality source for strong claim / outdated / marketing fluff.

****refuted=false**** ONLY if: claim is well-supported, current, and source quality matches claim strength.

Default to refuted=true if uncertain.

Structured output only. Evidence MUST be specific.

2. assistant (2026-06-06T17:37:42.801Z)

I'll verify this claim by examining the source paper and searching for corroborating or contradicting evidence.

3. assistant (2026-06-06T17:37:49.149Z)

I'll fetch the primary source paper to verify the claim is actually supported.

4. user (2026-06-06T17:37:56.113Z)

Analysis of Paper Content

I cannot provide a reliable answer to your questions because the provided content appears to be a corrupted or improperly encoded PDF file. The text consists primarily of binary/compressed data and garbled characters that cannot be meaningfully parsed to extract information about:

- The sport discussed
- Audio cue analysis methods
- Rally/exchange detection techniques
- Multimodal fusion approaches
- Specific performance metrics or datasets

****What I can confirm:**** The file header indicates this is a PDF document (`%PDF-1.2`), but the content streams are either compressed or corrupted, preventing extraction of the actual research findings.

****Recommendation:**** Please provide:

1. A properly decoded PDF file, or

2. The paper text in plain text/readable format, or
3. A direct link to the published paper

This will enable accurate analysis of the specific claims about audio discrimination, HMM fusion results, and the sport/dataset involved.

[Binary content (application/pdf, 188.7KB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\tool-results\webfetch-1780767080102-6svilu.pdf]

5. user (2026-06-06T17:37:57.934Z)

Web search results for query: "Kijak audio-visual HMM tennis video structure analysis rally ball hits applause"

Links: [{"title": "HMM BASED STRUCTURING OF TENNIS VIDEOS USING VISUAL AND AUDIO CUES", "url": "https://www.irisa.fr/metiss/gravier/biblio/03/kijak-icme03.pdf"}, {"title": "(PDF) A scheme for racquet sports video analysis with the combination of audio-visual information", "url": "https://www.researchgate.net/publication/238439853_A_scheme_for_racquet_sports_video_analysis_with_the_combination_of_audio-visual_information"}, {"title": "(PDF) Automatic video structuring based on HMMs and audio visual integration", "url": "https://www.academia.edu/13873617/Automatic_video_structuring_based_on_HMMs_and_audio_visual_integration"}, {"title": "Audiovisual Integration for Racquet Sports Video Retrieval | SpringerLink", "url": "https://link.springer.com/chapter/10.1007/11811305_74"}, {"title": "AUDIOVISUAL INTEGRATION FOR TENNIS BROADCAST STRUCTURING", "url": "https://www.irisa.fr/metiss/publis/kijak-cbmi03.pdf"}, {"title": "(PDF) Automatic annotation of tennis games: An integration of audio, vision, and learning", "url": "https://www.academia.edu/14129713/Automatic_annotation_of_tennis_games_An_integration_of_audio_vision_and_learning"}, {"title": "(PDF) Tennis video abstraction from audio and visual cues", "url": "https://www.researchgate.net/publication/4146865_Tennis_video_abstraction_from_audio_and_visual_cues"}, {"title": "Baseball and Tennis Video Annotation with Temporal Structure Decomposition", "url": "https://www.larry-lai.com/doc/MMSP08_Temporal.pdf"}]

I found relevant research on this topic. Here are the search results:

The search returns several academic papers related to Kijak's work on tennis video analysis using audio-visual information and HMMs (Hidden Markov Models):

****Key Research:****

The system analyzes tennis videos by extracting visual features like color and motion intensity, and audio features such as ball hits and applause, which are characteristics of rally scenes.

Using HMMs allows classification of tennis video into four defined units, achieving 93% accuracy in identifying point boundaries.

The audiovisual integration scheme is designed for racquet sports video structure indexing and highlight generation, combining visual and audio features rather than using low-level features alone.

The approach uses Hidden Markov Models for structure analysis of videos to merge audio and visual cues, validated in the domain of tennis videos.

****Key Applications:****

The detected auditory features including ball hitting and applause are used for identifying interesting events such as missed serves, aces, rallies and replays in tennis video.

The most relevant papers appear to be:

- "HMM Based Structuring of Tennis Videos Using Visual and Audio Cues" by Kijak et al.
- "Audiovisual Integration for Tennis Broadcast Structuring" by Kijak, Gravier, Oisel, and Gros (2003)

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. assistant (2026-06-06T17:38:02.451Z)

The PDF didn't parse via WebFetch. The search confirms the paper is about TENNIS (not badminton) and uses HMMs with audio-visual cues. The supporting quote mentions "ball hits" and "applause" which matches. Let me verify the specific claims more carefully by fetching the readable PDF locally and checking the companion paper.

7. user (2026-06-06T17:38:04.898Z)

pdfplumber available

8. user (2026-06-06T17:38:11.329Z)

Analysis of PDF Content

I cannot provide the requested information because the document provided is a compressed PDF file that appears to be corrupted or improperly formatted. The content displayed consists entirely of binary PDF encoding data rather than readable text.

To answer your questions about:

- The sport discussed
- Audio cues (ball hits, applause) and rally scene discrimination
- HMM fusion of audio-visual modalities
- Specific classification rates

I would need access to the actual, readable content of the paper.

****Recommendation:**** Please provide either:

1. A properly extracted text version of the document
2. A readable PDF that can be properly parsed
3. The relevant excerpts in text format

This would allow me to provide the specific numerical results and technical details you're seeking while adhering to proper citation practices.

[Binary content (application/pdf, 315.1KB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\42211f80-2169-4025-ad9d-99c1aa29ff70\tool-results\webfetch-1780767487469-2kpnam.pdf]

9. assistant (2026-06-06T17:38:13.778Z)

pdfplumber is available and both PDFs were saved locally. Let me extract the text directly from the primary source PDF.

10. user (2026-06-06T17:38:17.373Z)

=== PAGE 1 ===

HMMBASEDSTRUCTURINGOFTENNISVIDEOSUSINGVISUALANDAUDIOCUES

E.Kijak^{1,2},G.Gravier²,P.Gros²,L.Oisel¹ andF.Bimbot²

¹ ThomsonmultimediaR&D,France ² IRISA

1,av. deBelle-Fontaine CampusUniversitairedeBeaulieu

35510Cesson-Sevigne,France 35042RennesCedex,France

{ewa.kijak,lionel.oisel}@thomson.net {ggravier,pgras,bimbot}@irisa.fr

ABSTRACT are used to detect plays in the original flow. In a second

step, a detection of audience or commentator excitement

This paper focuses on the use of Hidden Markov Models

is used to select the most exciting plays [3]. The reverse

(HMMs)forstructureanalysisofvideos,anddemonstrates

schemehasalso beenproposedin[4]: cheeringsoundsare

how they canbeefficientlyappliedto mergeaudioandvi-

firstextractedtolocatesanimportantevent. Visualanalysis

sual cues. Our approach is validated in the particular do-

is then performed in cheering sounds regions to recognize

main of tennis videos. The basic temporal unit is the video event.

shot. Visual features are used to characterize the type of shot. In this paper, we propose a tennis scene classification method that automatically labels each segment of a video shot. The video structure parsing relies on the analysis of a tennis video with one of the following four types: first missed serve, rally, replay, and break. Hidden Markov models (HMMs) are used to merge audio-visual information, and typical tennis scenes are identified. In addition, each shot is assigned to a level in the hierarchy described in match. Domain knowledge is implicitly taken into account through the learning process associated with HMMs. This implicit approach is more suitable than heuristic rules oriented schemes because the learning process can derive more complete knowledge. In addition, previous works that use video structure parsing consists in extracting logical story units from the considered video. It's a mandatory step to efficiently organize and retrieve video contents. The structure to be estimated relies on the nature of the video. In our solution differs from existing methods in the following main points: (1) it combines visual/audio cues simultaneously; (2) different levels of classifications are generated over the entire tennis video at different levels; and (3), the segmented content are semantically meaningful to users. This paper is organized as follows. Section 2 gives an overview of the system and briefly describes the exploited broadcasted video data consist of several multimodal audio and visual features. Section 3 introduces the structure information streams. Numerous approaches presented in the literature have discussed how to extract high-level semantic events from video. To achieve such a goal, algorithms have to be dedicated to one particular type of videos such as sports or news. In this context, most of the approaches use the individual visual or audio cues. Only a few attempts dealing with multimodal analysis of audio-visual information have appeared recently, such as [1] on sports, and [2] on news. These approaches rely on a successive segmentation into shots by detecting cuts and dissolve transitions of visual and audio features. In a first step, visual features are extracted from the image stream, one

=== PAGE 2 ===

Domain Knowledge(cid:13)
 Sport game semantics(cid:13)
 &(cid:13) HMM Training(cid:13)
 Sport video syntax(cid:13)
 Shot Detection(cid:13) Shot Extraction(cid:13) Shot Classification(cid:13) Shot Parsing(cid:13) Program Syntax(cid:13)

States(cid:13)
Transitions(cid:13) Observations(cid:13)
Video program(cid:13)
frames(cid:13)

Fig.1. Structure analysis system overview

keyframe along with its image features (color and motion intensity), and from the audio stream, sound classes such as ball hits, applause, or speech. The segmented video results in an observation sequence which is parsed by a HMM process (Figure 1). As a result, each shot is assigned to a level in the hierarchy described in terms of point, game and close-up, and audience. In a tennis video production, global set, and is labelled with one of these four types: first missed views contain much of the pertinent information. The reserve, rally, replay, and break. maining information relies on the presence or the absence Such a tennis scene classification and segmentation is of non global views but is independent of the type of these possible because tennis video is characterized by a typical views. production style, known as the tennis video syntax.

2.2. Visual Features

2.1. Tennis Syntax

We used visual features to identify the global views within Sport video production is characterized by the use of all the extracted keyframes. The process can be divided into a number of cameras at almost fixed positions. Consider two steps. First, a keyframe K representative of a global view is selected without making any assumption about the relevant information is selected by the producer, and broadcast playing area color. Once K has been found, each keyframe is characterized by a similarity distance to K . In the following, we motivate the choice of the visual features we use, before briefly describing the identification process. by the camera filming the whole court is selected (we name Global views are characterized by a rather homogeneous global view in the following). After the end of the rally, color content (the colors of the court and its surrounding), the player who has just carried out an action of interest is although medium and audience views are characterized by captured with a close-up. As close-up views never appear scattered color content. In addition, a global view must during a rally but right after or before it, global views are capture each time the main part of the court, whereas in generally significant of a rally. Another example consists in close-up views, the camera is generally tracking the player. replay that are notified to the viewers by inserting special The first class can thus be characterized by a small camera motion while the second implies important camera transitions. We identify here four typical patterns in tennis videos. Based on these observations, we choose two features that are: first missed serve, rally, replay, and break. To identify global views: (1) a vector of dominant colors F break class includes the scenes unrelated to games, such as its spatial coherency C , and (2) activity A that reflects as commercials. A break is characterized by an important camera motion during a shot. succession of such scenes. It appears when players change The visual similarity measure $v(K_1, K_2)$ between two ends, generally every two games. We also take advantage keyframes K_1 and K_2 is defined as a

weighted function
of the well-defined and strong structure of tennis broadcast. of the spatial coherence,
the distance function between the
On the contrary of time-constrained sports that have a relatively dominant color vectors, and the activity:
tively loose structure, tennis is a score-constrained sport that $v(K1, K2) = (1)$
can be broken down into sets, games and points (Figure 2).
 $w1 |C1?C2| + w1 |d(F1, F2)| + w3 |A1?A2|$
The different types of views present in a tennis video
can be divided into four principal classes: global, medium, where $w1, w2$, and $w3$ are the weights.
=== PAGE 3 ===
To find a keyframe K representative of a global view, SET(cid:13) BREAK(cid:13)
ref
we consider dominant color ratios. In a reduced subset of MATCH(cid:13) S(cid:13) Br(cid:13) S(cid:13)
Br(cid:13) S(cid:13)
keyframes whose highest percentage of dominant color is
more than 50%, we apply the least median square method,
GAME(cid:13)
to minimize the median distance of all the remaining key- G(cid:13) Br(cid:13) G(cid:13)
G(cid:13) Br(cid:13) G(cid:13) G(cid:13) Br(cid:13) G(cid:13) G(cid:13)
frames to a randomly selected keyframe. This leads to ob-
tain K (see details in [7]). The visual similarity measure POINT(cid:13)
ref
 $v(K, K)$, denoted by v in the following, is then com- P(cid:13) P(cid:13) P(cid:13) P(cid:13)
P(cid:13)
t ref t
puted between each keyframe K and K .
t ref
First(cid:13)
2.3. Audio Features and $s, d, e, r, r, v, a, e, l, l$ (cid:13) y(cid:13)
Replay(cid:13)
Generally speaking, there are two approaches to combine
Rally(cid:13)
audio and visual features for structure analysis: one is to
combine audio and visual features into a single audiovisual
feature vector before the classification. The other consists
in classifying separately according to each modality before
Fig. 3. Content hierarchy of broadcast tennis video
integrating the classification results.
In this work, an intermediate strategy is used which con-
where $p(v | j)$ and $p(d | j)$ are given by a smoothed his-
sists in extracting separately ?high level? audio and visual t t
togram, and $P[a | j]$ is the product over each sound class
cues, the classification being made using the audio and vi- t
kof the discrete probability $P[a(k) | j]$.
sual cues simultaneously. As mentioned previously, the video t
Segmentation and classification of the whole observed
stream is segmented in a sequence of shots. Since shot
boundaries are more suitable for structure analysis than boun-
sequence $o = 0102...0$
T
into the different structural ele-
ments are performed simultaneously using a Viterbi algo-
daries extracted from the soundtrack, the shot is considered
as the base entity and features describing the audio content
rithm. The most likely state sequences = $s1...s$
T
is given
by
for each shot are used to provide additional information. For
each shot, a binary vector a describing which audio classes,
among speech, applause,
t
ball hits, noise and music, are
 $s? = \arg \min$
s

ax $\ln p(s) + (\text{cid}:88)\ln b$
st
(o
t
) (3)
t
present in the shot is extracted from an automatic segmentation of the audio stream. The automatic soundtrack segmentation is carried out using a Viterbi based system with a ten-mentation is carried out using a Viterbi based system with nis game, the four HMMs are connected to a higher-level Gaussian mixtures for each sound class and combination of HMM which represents the tennis syntax as illustrated in sound classes. The soundtrack is segmented independently Figure 3. The transition probabilities between states of the and the vectors a t are created from this independent seg- higher-level HMM result entirely from a priori information, mentation. while the transition probabilities of the sub-HMMs are estimated.

3. STRUCTURE ANALYSIS

4. EXPERIMENTAL RESULTS

Prior information is integrated by deriving syntactical basic elements from the tennis video syntax. We define four basic Experimental data are composed of about 2 hours of manually labelled tennis video. About half of the data is used to missed serves and rallies, the two others denote non-game train the HMM while the remaining half is reserved for the phases (breaks and replays). Each of these units is modelled tests. by a HMM. These HMMs rely on the temporal relationships Recall and precision rates are given in Table 1 for three experiments: visual features only, audio features only and Each HMM state models either a single shot or a dissolved transition between shots. The observation consists of three elements: the visual similarity between the shot t a sub-divided into two parts: using audio features manually annotated (denote by ?man.? in Table 1) to validate the a which characterizes the presence or absence of the pre-t model, and using audio features extracted from automatic determined audio events. More formally, the probability of segmentation (denote by ?segm.? in Table 1). The audio an observation o conditionally to state j is then given by t probabilities are estimated from the manually annotated audio features.

j t t t t

=== PAGE 4 ===

Visual Audio features Audio-visual
features man. segm. man. segm.
Segmentation 71% 68% 44% 78% 66%
accuracy

Classification precision recall precision recall precision recall precision recall
recall

First serve 67% 63% 92% 66% 66% 5% 87% 75% 71% 41%
Rallies 97% 54% 95% 63% 62% 28% 96% 83% 82% 65%
Replay 100% 80% 68% 38% 75% 15% 81% 88% 86% 82%
Break 51% 93% 91% 81% 30% 90% 94% 81% 68% 92%

Table 1.

Classification and segmentation results with visual features only, audio features only and both audio-visual features

We observed that one single stream cannot provide enough information. 5. CONCLUSION

ough information to identify different scenes. Rallies high precision rate (97%) and breaks recall rate (93%) show that in this paper, we have presented an audio-visual model for low level visual features are efficient to identify global views tennis video structure parsing. Tennis structure is modelled from others. However, these features are not sufficient to by an HMM that integrates a priori information about tennis content and editing rules. HMMs provide an efficient a first serve, as indicated by rallies low recall rate (54%). framework to merge audio and visual cues, and to perform This means also that many global views are missed using simultaneously classification and segmentation into scenes. only visual features. Replays detection relies essentially on Experiment shows that audio information improves the performance of video classification, when features are quite reliable. correspond to correctly detected dissolve transitions. Indeed errors from feature extraction spread over the entire analysis process. It is therefore necessary to enhance Classification using reliable audio features increases pre-audio segmentation and classification. cision rates for first missed serve. Comparable results are To extend our tennis video structure analysis, we also obtained for rallies detection. This suggests that audio features intend to include a more sophisticated audio-video fusion features are effective to describe rally scenes. Indeed, a rally is modelled by refining audio and video content analysis. essentially characterized by the presence of ball hits and applause which happen at the end of the exchange, although

6. REFERENCES

a missed first serve is only characterized by the presence of ball hits. Therefore, there are less confusion between [1] W. Hua, M. Han, and Y. Gong, "Baseball scene classification missed first serves and rallies. Breaks precision rate in using multimedia features," in Proc. of Int'l Conf. on Multimedia and Expo, 2002. the presence of music, the break state works in this case like [2] W. Qi, L. Gu, H. Jiang, X.-R. Chen, and H.-J. Zhang, "Integrated commercial detector, and avoids false detections. On the in visual, audio and text analysis for news video," in Proc. contrary, replays are not characterized by a representative of Int'l Conf. on Image Processing, 2000. audio content, and almost all replays are missed. [3] B. Li and I. Sezan, "Event detection and summarization in The approach using only audio features automatically american football broadcast video," in SPIE Storage and Retrieval for Media Databases, v. 4676, pp. 202-213, 2002. segmentation as described in 2.3 is 76.9%, and precision is 46.3%. Ball hits and music are well classified, whereas [4] K. Kim, J. Choi, N. Kim, and P. Kim, "Extracting semantic information from basketball video based on audio-visual features," in Proc. of Int'l Conf. on Image and Video Retrieval, automatic audio segmentations spread over the structuration vol. 2383, pp. 278-288, 2002. process and drag the performance down. Another source of errors is the mismatch between the manually segmented [5] J. Assfalg, M. Bertini, A. Del Bimbo, W. Nunziati, and P. Pala, audio features on which the HMM parameters were estimated, and the automatically generated audio features. An Proc. of Int'l Conf. on Multimedia and Expo, 2002. improvement was obtained by estimating the HMM parameters [6] P. Chang, M. Han, and Y. Gong, "Extract highlights from

eters using automatically generated audio features, but it did
baseball game video with hidden Markov models, in Proc. of
not outclass the video-only system. Int'l Conf. on Image Processing, 2002.
Nevertheless, integrating multiple cues from different
[7] E. Kijak, L. Oisel, and P. Gros, "Temporal structure analysis
media improves performance. In particular, ball hits detec- of broadcast tennis video using hidden
Markov models, in
tion improves classification rates for rallies and first serves. SPIE Storage and Retrieval for Media
Databases, v. 5021,
And music detection prevents from false detection of breaks. pp. 289-299, 2003.

11. assistant (2026-06-06T17:38:32.289Z)

I now have the full text of the primary source. Let me analyze the claim against the actual evidence carefully.

The claim has three parts:

1. Audio cues (ball hits + end-of-exchange applause) are discriminative for rally vs non-rally scenes
2. Fusing audio with visual cues in an HMM improves rally classification
3. ...over either modality alone

Let me check the actual Table 1 numbers for rallies:

****Rallies classification (precision / recall):****
- Visual only: 97% / 54%
- Audio (manual): 95% / 63%
- Audio (segm/automatic): 62% / 28%
- Audio-visual (manual): 96% / 83%
- Audio-visual (segm/automatic): 82% / 65%

Now let me scrutinize each part against the data.

****Part 1 (audio discriminative for rally):**** The paper says "audio features are effective to describe rally scenes" because rallies are characterized by ball hits + applause. Supported by the quote AND the data: audio-manual achieves 95%/63% for rallies. So Part 1 is supported, with the important caveat that this is true primarily with RELIABLE/manually-annotated audio. With AUTOMATIC audio segmentation, rally recall collapses to 28% (precision 62%), and the paper explicitly states the audio-only automatic system "gives poor performance."

****Part 2 + 3 (fusion in HMM improves over EITHER modality alone):**** This is the crux. Compare audio-visual vs the best single modality for RALLIES:
- Best single modality: visual 97%/54%, or audio-manual 95%/63%
- Audio-visual manual: 96%/83% ? recall jumps from 54%/63% to 83%, with precision ~96%. This clearly beats either modality alone on the combined precision+recall (especially recall).
- Audio-visual automatic (segm): 82%/65% ? beats visual-only recall (54% to 65%) and beats audio-only-automatic dramatically. Precision (82%) is below visual-only (97%) but recall and overall balance improves.

So for the RELIABLE-feature case, fusion clearly improves rally classification over either modality alone (recall 54/63 to 83, precision held at ~96). The paper's own conclusion: "audio information improve the performance of video classification, when features are quite reliable."

Now the critical skeptical checks:

12. user (2026-06-06T17:38:45.104Z)

Web search results for query: "tennis video segmentation audio ball hits applause HMM rally classification critique limitations later work"

Links: [{"title": "(PDF) Automatic video structuring based on HMMs and audio visual integration", "url": "https://www.academia.edu/13873617/Automatic_video_structuring_based_on_HMMs_and_audio_visual_integration"}, {"title": "HMM BASED STRUCTURING OF TENNIS VIDEOS USING VISUAL AND AUDIO CUES", "url": "https://www.irisa.fr/metiss/gravier/biblio/03/kijak-icme03.pdf"}, {"title": "Tennis video abstraction from audio and visual cues F. Coldefy \u0001 P. Bouthemy"}]

```
\u0001", "url": "https://www.irisa.fr/metiss/guig/biblio/04/coldefy-mm-sp-04.pdf"}, {"title": "(PDF) Classification of Tennis Video Types Based on Machine Learning Technology", "url": "https://www.researchgate.net/publication/352217162_Classification_of_Tennis_Video_Types_Based_on_Machine_Learning_Technology"}, {"title": "MULTIMODAL SEGMENTAL-BASED MODELING OF TENNIS VIDEO BROADCASTS", "url": "https://cecs.uci.edu/~papers/icme05/defevent/papers/cr1338.pdf"}, {"title": "(PDF) Tennis video abstraction from audio and visual cues", "url": "https://www.researchgate.net/publication/4146865_Tennis_video_abstraction_from_audio_and_visual_cues"}, {"title": "TennisExpert: Towards Expert-Level Analytical Sports Video Understanding", "url": "https://arxiv.org/pdf/2603.13397"}, {"title": "AUDIOVISUAL INTEGRATION FOR TENNIS BROADCAST STRUCTURING", "url": "https://www.irisa.fr/metiss/publis/kijak-cbmi03.pdf"}, {"title": "3D Convolutional Networks for Action Recognition: Application to Sport Gesture Recognition", "url": "https://arxiv.org/pdf/2204.08460"}]
```

Based on the search results, here's what I found about tennis video segmentation using audio, ball hits, applause, HMM, and rally classification:

Overview of the Approach

HMMs allow classification of tennis video into four defined units, achieving 93% accuracy in identifying point boundaries. Audio features include sound classes such as ball hits, applause, or speech.

System Features

A rally is essentially characterized by the presence of ball hits and applause which happen at the end of the exchange. Audio features include classifications for speech, applause, ball hits, noise, and music, represented as a binary vector for each shot.

Limitations of Single Modality

One single stream cannot provide enough information to identify different scenes. While low level visual features are not sufficient to distinguish a rally from a simple view of the court or from a first serve, which means many global views are missed using only visual features.

Multimodal Integration

The results show that combining audio and visual features is more effective. Rallies achieved a high precision rate (97%) and breaks a recall rate (93%) showing that low level visual features are efficient to identify global views from others.

Later Work

More recent approaches have moved beyond HMMs. Modern methods employ a lightweight, purely audio-based detection model that trains a neural network to identify the distinct acoustic signatures of racquet-ball impacts, augmented by overlaying various background sounds like crowd cheering, clapping, and announcer voices.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

13. assistant (2026-06-06T17:39:00.906Z)

The search corroborates the claim and confirms the source paper's findings are real. The claim is a faithful summary of the paper. Let me note one critical caveat for a skeptical assessment: this is a 2003 broadcast-tennis study, and the strong fusion result depends on RELIABLE audio features.

Let me do one final check on whether the claim is "outdated" or contradicted by later work, since the field has moved on, and whether the claim's specific wording is over-stated relative to the data.

Key observations for my verdict:

****Is the claim supported by the quote and the paper?**** Yes, strongly:

- The quote directly states audio features "are effective to describe rally scenes" and rallies are "characterized by the presence of ball hits and applause which happen at the end of the exchange."
- The quote states "integrating multiple cues from different media improves performance. In particular, ball hits detection improves classification rates for rallies."
- Table 1 backs this quantitatively: rally recall improves from 54% (visual-only) / 63% (audio-only-manual) to 83% (audio-visual-manual), precision held at ~96%. So fusion beats EITHER modality alone for rallies ? the claim's strongest assertion is directly evidenced.

****Counter-evidence / qualifications (the skeptical case):****

- The fusion improvement holds robustly only with RELIABLE/manually-annotated audio. With AUTOMATIC audio segmentation, audio-only rally recall collapses to 28% and the paper says audio-only-automatic "gives poor performance" and "did not outclass the video only system." However, audio-visual-automatic (82%/65%) still beats visual-only recall (54%), so fusion still helps even with automatic features. This is a qualification, not a refutation.
- It is a 2003 paper on BROADCAST TENNIS, not badminton, and not raw/local single-camera video. The "applause" cue depends on broadcast crowd audio. This limits direct transfer to the user's context but does NOT make the claim false as stated (the claim is about the general principle and cites this paper as evidence).
- Small dataset (~2 hours, half train/half test). But that's normal for the era and the claim isn't an extraordinary one.

The claim is accurately stated, faithfully reflects the primary source (which I read in full), and is corroborated by the paper's own quantitative table. It does not overreach: it says audio is "discriminative" (true) and fusion "improves... over either modality alone" (directly shown in Table 1 for rallies). No credible source contradicts it. The source is the primary, peer-reviewed paper (ICME 2003) ? appropriate quality for this modest, well-scoped claim.

This claim survives refutation.

14. user (2026-06-06T17:39:15.423Z)

Structured output provided successfully

15. assistant (2026-06-06T17:39:19.267Z)

I verified the claim against the full primary source and corroborating literature. Verdict: NOT REFUTED (high confidence). The claim faithfully reflects Kijak et al. 2003 (ICME), confirmed by the paper's own Table 1, with the bounded caveat that the strong fusion gain depends on reliable audio features and the study is on broadcast tennis rather than badminton.