

claude-windows / 56b7f849-5d56-4770-93e0-b3b4da943700

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\56b7f849-5d56-4770-93e0-b3b4da943700\subagents\workflows\wf_49729ae7-962\agent-a6daa390264b2e83c.jsonl`
- SHA-256: `3a787e76c44a77999e096739619b924276700edd856bf5faecb891ee5656bffa`
- Source modified: `2026-06-07T08:15:59+00:00`
- Imported at: `2026-07-05T16:48:27+00:00`
- project: `wf_49729ae7-962`
- session_id: `56b7f849-5d56-4770-93e0-b3b4da943700`
```

Transcript

1. user (2026-06-07T08:15:11.343Z)

You are a tough evaluator of how DEBUGGABLE an observability report is for a README-only reader.
Scenario: collector_missed_rally (tool: collector)
GROUND TRUTH (what actually went wrong): There is a long unlabeled gap (~26s) between rally 2 and rally 3, which may hide an omitted/missed rally. The reviewer should watch that span and add the rally if one is there. It is a review flag, not a blocker.
The report is expected to point the reader toward FAQ: "what-does-validation-check".

THE FRESH READER ANSWERED:

```
- diagnosis: The report shows a `possible_missed_rally` warning indicating a 26-second unlabeled gap between Rally 2 (ends at 14.0s) and Rally 3 (starts at 40.0s). According to the FAQ, this is a "clean omission" ? a real hole in the timeline where a rally was likely forgotten during annotation. The vendor's CVAT export is missing an entire rally in that time period.
- fix: Open the video at the 14.0s to 40.0s interval and verify whether a rally occurs there. If one exists, add it to the canonical CSV with the correct start_time, end_time, shots_count, ending_reason, and last_shot_outcome. Re-run `validate-delivery` to confirm the gap is resolved, then use the review sheet to document the correction before importing with `import-local`.
- faq_section used: How are *missed rallies* caught? (the worst error)
- could_self_debug: true
- confusing_parts: []
```

Also read the actual report for context:

C:\Users\avidu\AppData\Local\Temp\obseval\reports\collector_missed_rally.report.md

Judge strictly: did the reader land on the correct ROOT CAUSE (cause_correct)? Did they find/consult the expected FAQ (faq_correct)? How large is the remaining gap (none/minor/major)? Then propose ONE concrete, specific refinement to the REPORT wording, the spec rule message, or the FAQ that would have made this scenario unambiguous to self-debug ? or "none" if it is already excellent. Don't be generous; if anything was confusing, propose a fix.

2. assistant (2026-06-07T08:15:15.102Z)

I'll read the actual report to evaluate the reader's answer against ground truth.

3. user (2026-06-07T08:15:15.588Z)

```
1      # Sports Data Collector ? Delivery / Ingestion Report
2
3      ## Run
4
5      | Field | Value |
6      | --- | --- |
7      | Tool | sports-data-collector |
8      | Video id | pilot_gap (human_id) |
9      | Generated | 2026-06-07T08:10:51Z |
```

```

10 | Schema | 1.0 |
11
12 ## Video
13
14 | Field | Value |
15 | --- | --- |
16 | Duration | 10m00s |
17 | FPS | 59.94 fps (declared) |
18
19 ## Processing stages
20
21 **parse** ? `ok` (0.0s)
22   - rallies_parsed: 3
23   - _source_format_: cvat
24   - _fps_: 59.94
25   - _sampling_step_: None
26
27 **validation** ? `warn` (0.0s)
28   - issues_blocker: 0
29   - issues_high: 0
30   - issues_medium: 1
31   - issues_low: 0
32   - n_blockers: 0
33   - _gate_: PASS
34
35 ## Findings & issues
36
37   - [WARN] **possible_missed_rally** ? 26.0s of unlabeled time between Rally 2 (ends
14.0s) and Rally 3 (starts 40.0s) ? check for an unlabeled rally in the gap.
38     ? see **docs/INGESTION_PIPELINE.md#what-does-validation-check**
39     ? evidence: rally_number=3
40
41 ## Verdict
42
43 **PASS (with warnings)**
44

```

4. assistant (2026-06-07T08:15:20.032Z)

Now let me consult the FAQ that the report points to, to evaluate whether the reader landed on the correct root cause and whether their interpretation of "clean omission" matches the spec.

5. user (2026-06-07T08:15:20.286Z)

docs\INGESTION_PIPELINE.md

6. user (2026-06-07T08:15:22.818Z)

```

1      # Rally Ingestion Pipeline ? End-to-End FAQ
2
3      How a vendor's rally annotations become training/eval ground truth, and how we
4      validate each delivery against our conventions. This is the operational companion
5      to the annotation spec (`docs/annotation/ROORA_BRIEF.md` + `ANNOTATION_GUIDE.md`).
6
7      ```
8      vendor delivery                                OUR side (collector)                                modeling
(indexer)
9      ?????????????????? convert ?????????????????? validate ??????????????????
10     ? CVAT export ? ?????????????? ? canonical CSV ? ?????????????? ? report .md ?
11     ? (boxes + attrs)? ? (per rally, ? ? report .json ?
12     ? OR ? ? ? decimal secs)? ? review .csv ???? human
verifies
13     ? canonical CSV ? ?????????????????? ??????????????????
flagged rallies

```

```

14      ??????????????????      ?      ?
15      ? import-local (validated) ?
16      ?      golden labels
17      ?????????????????? export ??????????????????
18      ? SQLite (rally ? ?????????? ? <id>.csv      ?
start,end (indexer GT)
19      ? tables)      ?      ? manifest.json?
20      ??????????????????      ??????????????????
21      ?
22      slice clips / calibrate /
train
23      ...
24
25      ## The four commands (in order)
26
27      ```bash
28      # 1. Normalise a CVAT export ? our canonical per-rally CSV (skip if the vendor sent a
CSV)
29      python -m src.cli.curate convert-cvat --input job.xml \
30      --video-path match.mp4 --out canonical.csv --issues issues.txt
31
32      # 2. Validate the delivery ? report + annotatable review sheet (the QA gate)
33      python -m src.cli.curate validate-delivery --input job.xml \
34      --video-path match.mp4 --video-id pilot_badminton_01 --out-dir data/validation
35      # ...or validate a CSV directly:
36      python -m src.cli.curate validate-delivery --csv canonical.csv --video-id
pilot_badminton_01
37
38      # 3. After human review, import the (corrected) CSV into the DB
39      python -m src.cli.curate import-local --sport badminton \
40      --video-path match.mp4 --annotations-path canonical.csv --fps 59.94
41
42      # 4. Export indexer-ready ground truth (start,end CSVs + manifest)
43      python -m src.cli.curate export --out-dir data/eval_export
44      ...
45
46      ---
47
48      ## FAQ
49
50      ### What delivery formats do we accept?
51      Both. A **CVAT export** (image- or track-mode XML ? rally fields carried as box
52      attributes) is normalised by `convert-cvat`; an **already-canonical CSV** is taken
53      as-is. Vendors may keep working in CVAT ? we do **not** require them to hand-author
54      CSVs.
55
56      ### What is the "canonical" CSV?
57      One row per rally:
58      `rally_number, start_mmss, end_mmss, start_time, end_time, shots_count, ending_reason,
last_shot_outcome, score_before, score_after, notes`.
59      `start_time`/`end_time` are **decimal seconds** ; `import-local` reads those.
60
61      ### How are timestamps handled? (the fps thing)
62      We recompute each rally's interval from its **frame range × the true fps**
63      (`seconds = frame / fps`). A vendor's CVAT time attributes are **ignored** ? a
64      step-sampled job that assumes 30 fps writes systematically wrong times. fps comes
65      from `--fps` or ffprobe of `--video-path`. `mm:ss` is also accepted everywhere.
66
67      ### What does validation check?
68      | code | severity | meaning |
69      |---|---|---|
70      | `overlap` | **blocker** | two rallies' intervals overlap (`end[N] > start[N+1]`) |
71      | `zero_duration` | **blocker** | `end <= start` (single-sample / stray box) |
72      | `frame_out_of_range` | high | a box frame is past the video's last frame |
73      | `shot_density` | medium | long rally with implausibly few shots ? likely an **inflated

```

```

span** (often hides a merged rally) |
74 | `possible_missed_rally` | medium | a long unlabeled gap between two rallies ? check
for an **omitted rally** |
75 | `ending_reason_vocab` | medium | `ending_reason` outside the agreed set |
76 | `outcome_vocab` | low | `last_shot_outcome` outside `{in,net,out,n_a}` |
77
78 ### Blocker vs. review ? what's the gate?
79 `validate-delivery` exits **non-zero if any blocker exists** (so it can gate a bulk
80 loop / CI). Non-blocker issues don't stop the gate but populate the **review sheet**
81 for the human pass. `import-local` independently rejects overlaps/zero-duration, so
82 bad ground truth can't slip into the DB.
83
84 ### How are *missed rallies* caught? (the worst error)
85 Two complementary checks, because a miss shows up two ways:
86 1. **Clean omission** (a real hole in the timeline) ? `possible_missed_rally` (gap
heuristic).
87 2. **Merged/inflated** (the missed rally got swallowed into a neighbour's box, leaving
88 no gap) ? `shot_density` flags the over-long, under-shot rally, and the human finds
89 the second rally inside it.
90 Neither is a guarantee ? they **surface regions to review**, they don't certify
91 completeness. (A future enhancement: a motion second-opinion over the raw video.)
92
93 ### What's the review sheet?
94 `

```

136 two runs to see what changed between deliveries.
137 - `**`<id>.report.md`**` ? the `**technical**` report: parse / validation / registration
138 stages plus `**findings**` (one per delivery issue, each citing this FAQ). Read this
139 to debug a delivery.
140 - `**`<id>.summary.md`**` ? a `**delivery summary**` for the vendor / data team: the
141 rally count and coverage plus the plain-language list of rallies to review ? no
142 internal codes or severities.
143 This is in addition to the existing ``<id>_report.md`` / ``_report.json`` / ``_review.csv``
144 (the human review sheet), which are unchanged.
145

7. assistant (2026-06-07T08:15:46.585Z)

Now I have full context. Let me evaluate carefully.

Ground truth: Long unlabeled gap (~26s) between rally 2 and rally 3 that may hide an omitted/missed rally. Reviewer should watch that span and add a rally if present. It's a review flag, not a blocker. Expected FAQ: "what-does-validation-check".

Reader's answer assessment:

1. **Root cause (diagnosis):** The reader correctly identified the ``possible_missed_rally`` warning, the 26s gap between Rally 2 (ends 14.0s) and Rally 3 (starts 40.0s), and correctly concluded a rally was likely forgotten/omitted in annotation. This matches ground truth. The core diagnosis is correct.

One inaccuracy: the reader asserted "The vendor's CVAT export is missing an entire rally in that time period" as a *fact*, but the FAQ explicitly says these checks "surface regions to review, they don't certify completeness" and "Neither is a guarantee." The reader over-stated a heuristic flag as a certainty. But the prescribed action (go look) is right.

2. **FAQ:** Expected FAQ is "what-does-validation-check" (the report's citation anchor ``#what-does-validation-check``). The reader cited "How are *missed rallies* caught? (the worst error)" instead. That is a *different* FAQ section. So strictly, they did not consult the exact expected section ? though the section they used is highly adjacent and arguably more directly explanatory of the root cause. The "what does validation check" table does contain the ``possible_missed_rally`` row. The reader landed on a sibling section. This is a partial FAQ miss.

3. **Fix:** Correct and complete ? watch 14.0?40.0s, add the rally with proper fields if present, re-run validate-delivery, document in review sheet, then import-local. Aligns with the pipeline.

4. **Severity/gap framing:** The reader did NOT explicitly call out that this is medium/review-flag-not-a-blocker (verdict is PASS-with-warnings). They treated it somewhat as "the worst error" (echoing the FAQ heading) and said the export "is missing an entire rally" ? slightly over-escalating a non-blocking review flag. Ground truth stresses "It is a review flag, not a blocker." The reader did note re-running to "confirm the gap is resolved," implying review-ability.

Let me assess the gap size and the FAQ-correct boolean.

- `cause_correct`: TRUE ? they nailed the gap, the interval, and the omission interpretation; action is right.
- `faq_correct`: The expected anchor is ``#what-does-validation-check``. The reader cited the "How are missed rallies caught?" section. Strictly, they cited a different section. However the report's finding line cites the validation-check anchor directly. They could self-debug regardless. Judging strictly: they did NOT consult the expected FAQ section ? FALSE. They used an adjacent one that happens to cover the same case, but the question asks if they found/consulted the *expected* FAQ.
- `gap`: minor ? diagnosis and fix correct, but two wording-level slips: (a) over-asserting "missing an entire rally" as fact vs. heuristic flag, (b) framing it as "the worst error" rather than a medium review flag.

Now a concrete refinement. The strongest single fix: the report's finding line says "check for an unlabeled rally in the gap" but does not convey that this is a *heuristic* / not a *guarantee*

nor that it is a non-blocking review flag ? which is exactly what nudged the reader to over-assert certainty and reach for the "worst error" framing. Also, the report cites `#what-does-validation-check` but the most explanatory text is in "How are missed rallies caught?".

The most useful, specific refinement is to make the report finding line state the severity/non-guarantee explicitly and point to the more diagnostic anchor.

8. user (2026-06-07T08:15:53.798Z)

Structured output provided successfully

9. assistant (2026-06-07T08:15:59.639Z)

I've completed the evaluation and submitted the structured output.

Summary of my judgment:

- **cause_correct: true** ? The reader correctly identified the root cause (26s unlabeled gap between Rally 2 and Rally 3 likely hiding an omitted rally) and prescribed the right remediation.
- **faq_correct: false** ? The report cites `#what-does-validation-check`, but the reader consulted/cited the adjacent "How are missed rallies caught?" section instead. Adjacent and explanatory, but not the expected anchor.
- **gap: minor** ? Two wording slips, both stemming from the report not encoding severity/uncertainty: the reader over-asserted "the export is missing an entire rally" as fact (the FAQ stresses these flags are not a guarantee), and over-escalated it as "the worst error" rather than a non-blocking medium review flag.
- **Refinement** ? Make the report finding line carry the severity, the heuristic/not-a-certainty caveat, and the non-blocker status explicitly, which would have prevented both slips.