

claude-windows / 45ea3a91-654d-4972-9cd3-65f35db54caf

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\45ea3a91-654d-4972-9cd3-65f35db54caf\subagents\workflows\wf_2a64bc88-2a8\agent-a35d1ecbaa0be395f.jsonl`  
- SHA-256: `f043386ee005f16b76a171879f8f44ed73239b9a2af559f1f7e7cbb62a431013`  
- Source modified: `2026-06-10T06:09:28+00:00`  
- Imported at: `2026-07-05T16:48:27+00:00`  
- project: `wf_2a64bc88-2a8`  
- session_id: `45ea3a91-654d-4972-9cd3-65f35db54caf`
```

Transcript

1. user (2026-06-10T05:59:28.045Z)

CONTEXT ? two sibling repos on a Windows machine:

- INDEXER: C:/Users/avidu/Projects/badminton-highlight-indexer ? local AI badminton(->racket-sports) highlight indexer + rally detector. FastAPI + SQLite + ffmpeg + GPU CV via WSL; Gemini = paid cloud AI. docs/ tree is rich. Current branch docs/next-steps-multisport (master = main branch).

- COLLECTOR: C:/Users/avidu/Projects/sports-data-collector ? golden-label acquisition/curation/ingestion CLI (system-of-record candidate for training corpus).

Project state (June 2026): scaffolding complete; owned-model roadmap agreed (docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md) but NO platform/owned-model implementation started; next committed platform slice = async job model. Licensing guardrail: MIT/Apache-2.0/BSD only for code+data+weights; paid LLMs are inference-only, never training-data sources. ENVIRONMENT RULES: this dev environment has NO GPU/torch/cv2 ? do NOT attempt to run the video pipeline or import torch/cv2 modules. Test suites are offline/fully-mocked and SAFE to run. You are READ-ONLY: never edit/create/delete repo files; running tests and git read commands is allowed.

QUALITY BAR: cite file paths (and line numbers where possible) for every finding. Report what IS in the code/docs, not what you assume. Be skeptical and concrete. Your final structured output is consumed programmatically by an orchestrator ? return raw data, not prose for a human.

TASK: full audit (engineering + testing + reliability + docs) of the COLLECTOR repo at C:/Users/avidu/Projects/sports-data-collector.

Steps: read README.md and all of docs/ (DECISIONS.md, PIPELINE_ROLE.md, INGESTION_PIPELINE.md, MANUAL_INGESTION_POSITIONING.md, annotation/*). Map the src/ tree (Glob src/**/*.py) and deep-read the core: src/cli/curate.py, src/cli/normalize_annotations.py, src/parsers/ (esp. cvat/rally_cvat.py), the db layer, models/validation, reporting/. RUN the test suite (pytest from repo root ? offline, safe); capture real counts.

ASSESS: (1) engineering quality and architecture vs its stated role as the SYSTEM-OF-RECORD for the training corpus (docs/MANUAL_INGESTION_POSITIONING.md) ? is it actually ready for that role? what's missing (provenance, maturity ladder, labeled_window persistence ? check which of these exist TODAY vs are documented as P1/P2); (2) known landmines: import-local DELETE-then-INSERT destructiveness, int(float()) rally_number PK behavior, pydantic abort-whole-import semantics, the normalize_annotations guard coverage; (3) the indexer->collector CONTRACT: does "curate export" output (manifest.json + per-video start,end CSV) actually match what the indexer's backend/eval/batch.py + run_phase3_eval.py consume (read those indexer files to cross-check field names/paths/Windows path normalization)? does convert-cvat's ATTR_MAP handle the real Roora CVAT attribute names (spaces, "(A-B)")? (4) test coverage gaps and quality; (5) reliability: what happens on partial/corrupt CVAT zips, duplicate imports, fps mismatch. Severity-rate as if the training corpus' integrity is on the line (it is ? this feeds the owned model).

2. user (2026-06-10T05:59:37.495Z)

```
README.md  
__pycache__
```

cleanup.py
config
data
docs
pyproject.toml
src
tests
DECISIONS.md
INGESTION_PIPELINE.md
MANUAL_INGESTION_POSITIONING.md
PIPELINE_ROLE.md
annotation

3. user (2026-06-10T05:59:37.649Z)

No files found

4. user (2026-06-10T05:59:37.649Z)

No files found

5. user (2026-06-10T05:59:51.044Z)

C:/Users/avidu/Projects/sports-data-collector/src/cli/curate.py
C:/Users/avidu/Projects/sports-data-collector/src/cli/ingest.py
C:/Users/avidu/Projects/sports-data-collector/src/cli/normalize_annotations.py
C:/Users/avidu/Projects/sports-data-collector/src/cli/__pycache__/curate.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/cli/__pycache__/ingest.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/cli/__pycache__/normalize_annotations.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/config.py
C:/Users/avidu/Projects/sports-data-collector/src/core/db.py
C:/Users/avidu/Projects/sports-data-collector/src/core/downloader.py
C:/Users/avidu/Projects/sports-data-collector/src/core/parsing.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/badminton.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/base.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/cricket.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/pickleball.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/tabletennis.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/tennis.py
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/__pycache__/badminton.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/__pycache__/base.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/__pycache__/cricket.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/__pycache__/pickleball.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/__pycache__/tabletennis.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/__pycache__/tennis.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/validator.py
C:/Users/avidu/Projects/sports-data-collector/src/core/__pycache__/config.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/__pycache__/db.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/__pycache__/downloader.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/__pycache__/parsing.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/core/__pycache__/validator.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/crawlers/badminton_crawler.py
C:/Users/avidu/Projects/sports-data-collector/src/crawlers/__pycache__/badminton_crawler.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/parsers/badminton/coachai.py
C:/Users/avidu/Projects/sports-data-collector/src/parsers/badminton/__pycache__/coachai.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/parsers/base_parser.py

C:/Users/avidu/Projects/sports-data-collector/src/parsers/cvat/rally_cvat.py
C:/Users/avidu/Projects/sports-data-collector/src/parsers/cvat/__init__.py
C:/Users/avidu/Projects/sports-data-collector/src/parsers/cvat/__pycache__/rally_cvat.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/parsers/cvat/__pycache__/__init__.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/parsers/__pycache__/base_parser.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/reporting/adapters.py
C:/Users/avidu/Projects/sports-data-collector/src/reporting/spec.yaml
C:/Users/avidu/Projects/sports-data-collector/src/reporting/__init__.py
C:/Users/avidu/Projects/sports-data-collector/src/reporting/__pycache__/adapters.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/src/reporting/__pycache__/__init__.cpython-313.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/test_cli.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_config.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_cvat_parser.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_db.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_downloader.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_local_import.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_normalize_annotations.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_parser.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_parsing.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_reporting.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_reporting_faq.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_schemas.py
C:/Users/avidu/Projects/sports-data-collector/tests/test_validator.py
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_cli.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_config.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_cvat_parser.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_db.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_downloader.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_local_import.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_normalize_annotations.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_parser.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_parsing.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_reporting.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_reporting_faq.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_schemas.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/tests/__pycache__/test_validator.cpython-313-pytest-9.0.3.pyc
C:/Users/avidu/Projects/sports-data-collector/config/config.yaml
C:/Users/avidu/Projects/sports-data-collector/docs/annotation/ANNOTATION_GUIDE.md
C:/Users/avidu/Projects/sports-data-collector/docs/annotation/INTAKE_MANIFEST_TEMPLATE.csv
C:/Users/avidu/Projects/sports-data-collector/docs/annotation/NDA_TEMPLATE.md
C:/Users/avidu/Projects/sports-data-collector/docs/annotation/ROORA_BRIEF.md
C:/Users/avidu/Projects/sports-data-collector/docs/DECISIONS.md
C:/Users/avidu/Projects/sports-data-collector/docs/INGESTION_PIPELINE.md
C:/Users/avidu/Projects/sports-data-collector/docs/MANUAL_INGESTION_POSITIONING.md
C:/Users/avidu/Projects/sports-data-collector/docs/PIPELINE_ROLE.md

6. user (2026-06-10T05:59:55.905Z)

1 # Sports Data Collector & Curator

```

2
3     A clean, extensible Python codebase and CLI toolchain designed to discover, ingest,
validate, and curate sports video datasets. The repository stores annotations and video metadata
in an SQLite database, supports sport-specific tables and schemas, manages video storage outside
of Git, and enforces strict commercial license compliance.
4
5     > **Pipeline role:** this repo is the *data producer* feeding the **badminton-highlight-
indexer**, which trains/evaluates an offline rally segmenter on this curated, commercially-
licensed data. See [docs/PIPELINE_ROLE.md](docs/PIPELINE_ROLE.md) for why `export`/`fetch` exist
and how the two repos fit together.
6
7     ---
8
9     ## Features
10
11     - **SQLite Backend**: Centralized `sports_metadata.db` with a generic `videos` table and
sport-specific annotation tables (`badminton_rallies`, `tennis_points`, `cricket_deliveries`,
`pickleball_rallies`, `tabletennis_rallies`).
12     - **Commercial License Flagging**: Each ingested source license is checked against a
configured whitelist (`MIT`, `Apache-2.0`, `BSD`, `CC0`, `CC-BY`, `Public-Domain`,
`Proprietary`) and the resulting `is_commercial` flag is stored alongside the video. Non-
commercial sources (like `CC-BY-NC-SA`) are flagged as non-commercial but still ingested ?
nothing is discarded automatically. The whitelist is read from `config/config.yaml` consistently
across the parser, validator, and CLI.
13     - **Overlap & Duration Validation**: Ensures chronological event bounds (e.g. rally
start/end times) do not overlap and that durations are positive.
14     - **Stroke-level CSV Parser**: Ingests ShuttleSet/CoachAI-style stroke CSV logs and
aggregates them into rally durations using configurable start/end time buffers.
15     - **Local Ingestion Mode**: Manually registers local videos and associated CSV/JSON
annotation files directly into the database as curated entries.
16     - **Interactive Curation CLI & Visual Play Harness**: Let's you query DB status, list
staging entries, review sample timestamps, and approve/reject them. It includes a built-in play
function that opens the video in your default browser at the exact rally segment (using HTML5
media fragments `#t=start,end` for local files or YouTube start time markers `&t=start` for
remote videos) for manual visual verification.
17
18     ---
19
20     ## Installation
21
22     Ensure you have Python 3.9+ installed.
23
24     1. Clone or download the repository.
25     2. Install the package and dependencies:
26         ```bash
27         pip install -e .           # core: curate / import-local / export / status
28         pip install -e ".[fetch]"  # adds yt-dlp for video acquisition (fetch / crawler)
29         pip install -e ".[dev]"    # test tooling (pytest, pytest-cov)
30         ```
31
32     ### External requirements (not pip-installable)
33
34     The `fetch` command and the ShuttleSet crawler shell out to external binaries.
35     They degrade gracefully (a clear error / skipped download) when missing, but are
36     **required** for video acquisition:
37
38     - **FFmpeg / ffprobe** ? used to probe the real resolution/fps of downloaded files.
39     - **A JS runtime (e.g. [deno](https://deno.com))** ? needed by yt-dlp to solve
40       YouTube's challenge for high-resolution formats. `fetch` auto-detects a deno
41       install at `~/deno/bin`.
42
43     Core curation/import/export/status commands work without any of these.
44
45     ---
46

```

```

47     ## Usage
48
49     All commands should be executed from the root of the project. Make sure `PYTHONPATH` is
set to the current directory:
50
51     ```bash
52     # On Windows PowerShell:
53     $env:PYTHONPATH="."
54     ```
55
56     ### 1. View Database Status
57     To check the current statistics of staging vs. curated items:
58     ```bash
59     python -m src.cli.curate status
60     ```
61
62     ### 2. Ingest External Datasets to Staging
63     To parse external raw stroke-level datasets (e.g. CoachAI format) and register them in
staging:
64     ```bash
65     python -m src.cli.ingest --sport badminton --parser coachai --file-path
data/staging_coachai_match.csv --license CC-BY-NC-SA --video-id staging_match_non_comm
66     ```
67
68     ### 3. Launch the Interactive Curation Tool & Play Harness
69     To review staging items, inspect their timestamps, play individual rallies in the
browser to visually verify their start and end bounds, and approve or reject them:
70     ```bash
71     python -m src.cli.curate curate
72     ```
73
74     ### 4. Manually Import Local Videos & Annotations
75     To register a local video file alongside its JSON/CSV annotation file (bypasses
staging):
76     ```bash
77     python -m src.cli.curate import-local --sport badminton --video-path
"data/videos/my_custom_match.mp4" --annotations-path "data/mock_local_badminton_anno.json"
--match-name "My Custom Local Match" --license Proprietary
78     ```
79     `import-local` works for every configured sport ? `badminton`, `tennis`, `cricket`,
`pickleball`, and `tabletennis` ? from either a JSON or a CSV annotations file. For the
canonical rally CSV schema and the rater-facing instructions used to produce these files (e.g.
via an annotation vendor), see [docs/annotation/](docs/annotation/).
80
81     ### 5. Fetch / Upgrade Videos to Best Resolution
82     Videos are always downloaded at the best resolution available** (via yt-dlp `bv*+ba/b
-S res,...`, merged to MP4); the actual width/height/fps is ffprobed and stored ? never assumed.
The `fetch` command (re)downloads videos for entries already in the DB and updates
`local_path`/`resolution`/`fps` in place, preserving all rally annotations and curation
status (no clean slate needed):
83     ```bash
84     # Upgrade every sub-720p video to best resolution, using a logged-in browser's cookies:
85     python -m src.cli.curate fetch --cookies-from-browser firefox
86
87     # A single video, forcing a re-download even if a good local copy exists:
88     python -m src.cli.curate fetch --video-id shuttleset_1 --force --cookies-from-browser
firefox
89     ```
90     Flags: `--sport`, `--video-id`, `--status {staging,curated,all}` (default `all`),
`--min-height` (skip videos already at/above it unless `--force`; default `720`), `--max-height`
(cap resolution, e.g. `1080`; default best available), `--force`, `--include-noncommercial`,
`--cookies-from-browser`/`--cookies`, `--player-client`. A global resolution cap can also be set
via `video_storage.max_height` in `config/config.yaml`. Re-runs are cheap and resumable; a
download that falls below `--min-height` is treated as a failure (discarded, DB left
unchanged) so a flaky low-res result never overwrites a good entry.

```

```

91
92     > **Authenticated cookies are strongly recommended.** Anonymous YouTube downloads are
throttled and non-deterministic (often falling back to 144p). Pass --cookies-from-browser
firefox` (or `chrome`/`edge`/`brave`, fully closed) ? or --cookies cookies.txt` ? to use a
logged-in (ideally Premium) session for fast, reliable best-resolution downloads. Chrome's app-
bound cookie encryption may block direct reads on Windows; Firefox reads cleanly even while
open.
93     >
94     > **JS runtime:** yt-dlp needs one (e.g. [deno](https://deno.com)) to solve YouTube's
challenges for high-res formats; `fetch` auto-detects a deno install at ~/deno/bin`. The
default --player-client` avoids `web`/`web_safari`, which YouTube forces into SABR streaming
(drops to 144p).
95
96     ### 6. Export Annotations as Indexer-Ready Eval/Training Sets
97     To turn curated annotations into ground-truth files the badminton-highlight-indexer can
score and train against, export per-video `start,end` CSVs plus a `manifest.json`:
98     ```bash
99     # Curated + commercially-safe videos across all supported sports (default):
100     python -m src.cli.curate export --out-dir data/eval_export
101
102     # One sport, also including staging entries:
103     python -m src.cli.curate export --sport badminton --include-staging
104     ```
105     The exporter writes <out-dir>/<sport>/<video_id>.csv` (decimal-second `start,end` rows)
and a `manifest.json` pairing each CSV with its video `local_path`/`video_url`, fps, resolution,
and license. The indexer keys videos by **file MD5**, not by our `video_id`, so the downstream
eval/training run must process the exact video file named in the manifest. Flags: --sport`
(restrict to one sport), --include-staging` (default: curated only), --include-noncommercial`
(default: commercial-safe only).
106
107     ---
108
109     ## Configuration
110
111     Adjust settings in `config/config.yaml`:
112     - **`database_path`**: Path to the SQLite database.
113     - **`video_storage`**: Local folder path and automatic download size limit.
114     - **`commercial_licenses`**: Whitelist of licenses allowed for commercial use.
115     - **`supported_sports`**: Allowed sports categories.
116
117     ---
118
119     ## Documentation
120
121     - [docs/PIPELINE_ROLE.md](docs/PIPELINE_ROLE.md) ? how this repo feeds the **badminton-
highlight-indexer** (`export` / `fetch`).
122     - [docs/DECISIONS.md](docs/DECISIONS.md) ? collector-local decision records
(**ADR-001**: keep explicit per-sport code instead of a registry, with an "adding a new sport"
checklist).
123     - [docs/annotation/](docs/annotation/) ? the vendor annotation package used to produce
golden rally labels:
124     - [ROORA_BRIEF.md](docs/annotation/ROORA_BRIEF.md) ? vendor brief + the canonical
rally CSV schema.
125     - [ANNOTATION_GUIDE.md](docs/annotation/ANNOTATION_GUIDE.md) ? the illustrated rater
guide (singles & doubles).
126     - [INTAKE_MANIFEST_TEMPLATE.csv](docs/annotation/INTAKE_MANIFEST_TEMPLATE.csv) ? per-
video intake sheet.
127     - [NDA_TEMPLATE.md](docs/annotation/NDA_TEMPLATE.md) ? mutual NDA template (have
counsel review).
128
129     ---
130
131     ## Running Tests
132
133     We use `pytest` for unit and integration testing:

```

```

134  ```bash
135  python -m pytest tests/
136  python -m pytest tests/ --cov=src --cov-report=term-missing # with coverage (~87%)
137  ```
138

```

7. user (2026-06-10T05:59:56.854Z)

```

1      # Positioning: manual (rally-annotator) golden-label ingestion
2
3      How `sports-data-collector` should be positioned to ingest manual rally labels
4      produced by the VLC [rally-annotator](https://github.com/avidullu/rally-annotator),
5      alongside vendor (Roora/CVAT) deliveries, and serve them to the
6      badminton-highlight-indexer eval/training pipeline.
7
8      > Source: a 10-agent analysis workflow across this repo, the indexer, and the
9      > annotator, verified against the real `Adarsh and Avi.rallies.csv`. Confidence: high.
10
11     ## Bottom line
12
13     Position the collector as the single canonical system-of-record for golden rally
14     labels: the one place where any annotation source ? manual VLC lean CSV, the rich
15     review sheet, or vendor CVAT/CSV ? becomes a validated, license-gated, registered rally
16     record. The indexer consumes only the collector's narrow export contract:
17     `manifest.json` + a per-video `start,end` CSV (`curate.py::export_eval_set`;
18     `calibrate_wasb.py::load_golden_gts` reads `start_time`/`end_time` only).
19
20     The load-bearing property to protect: the seam to the indexer stays intervals-only.
21     Everything else ? `ending_reason`, `shots_count`, scores, provenance, maturity, the
22     labeled window ? is validated at the hub and rides in the manifest, never the eval
23     CSV. A new sport or source becomes an adapter change at the hub, never an indexer
24     change.
25
26     ## Role in the pipeline
27
28     ```
29     producers ?????????????? collector (hub) ?????????????? consumer
30     ? VLC rally-annotator      parse ? validate ? license-gate  badminton-highlight-
31     (lean 5-col CSV)           ? register (SQLite) ? export      indexer (eval/training)
32     ? manual review sheet                                           reads ONLY:
33     (rich 14-col CSV)         everything but (start,end) is      manifest.json +
34     ? vendor Roora / CVAT     validated then dropped before    per-video start,end
35
36     (CVAT XML / canonical)    the indexer sees it
37     ```
38
39     ## It already works today (zero code changes)
40
41     The lean annotator CSV header `rally_number,start_time,end_time,ending_reason,sport`
42     matches the `import-local` `DictReader` exactly (`curate.py::_parse_local_annotations`),
43     and the validator accepts gapped/partial coverage as-is
44     (`validator.py::validate_badminton_rallies` checks only positive duration + no overlap).
45     Verified end-to-end on the real file ? 40 rallies ? export ? 40-interval indexer CSV.
46
47     ```bash
48     python -m src.cli.curate import-local --sport badminton \
49     --video-path "?/Adarsh and Avi.MP4" \
50     --annotations-path "?/Adarsh and Avi.rallies.csv" --license Proprietary --fps 59.94
51     python -m src.cli.curate export # ? manifest.json + <video_id>.csv (start,end)
52     ```
53
54     ## P0 ? the two-landmine guard (`normalize_annotations.py`, shipped)
55
56     Partial human data has two failure modes the raw path handles badly. The
57     `src/cli/normalize_annotations.py` pre-import normalizer fixes both, as a standalone

```

```

56     step** that feeds `import-local` (it does **not** touch the DB schema or the
57     `rally_number`-ordered export):
58
59     1. **An unusable time row aborts the *entire* import.** A blank/missing time makes
60         `safe_float_required` raise (`parsing.py`); a reversed `end_time <= start_time` or a
61         negative `start_time` makes the `BadmintonRally` pydantic validator raise
62         (`schemas/badminton.py::validate_timestamps`). Either way the `try/except` in
63         `import-local` turns it into `sys.exit(1)`, killing the whole video. ? the normalizer
64         **drops** such rows (mirroring the indexer loader, which already skips `not end >
65         start`).
66     2. **A fractional/duplicate `rally_number` (e.g. `7.5`, `16.5`) silently corrupts
67         data.**
68         `int(float("7.5"))` ? 7` (`parsing.py::safe_int`) collides on the
69         `(video_id, rally_number)` primary key, and the DELETE-then-INSERT upsert
70         (`db.py::insert_badminton_rallies`) destroys the colliding rally. ? the normalizer
71         **renumbers `1..N` by `start_time`, recording the original id for traceability.
72
73     It also converts `mm:ss`/`HH:MM:SS` ? decimal seconds and folds a pipe-delimited
74     `ending_reason` ("forced_error|other") to its primary token.
75
76     ```bash
77     python -m src.cli.normalize_annotations --in "Match.rallies.csv" --out clean.csv
78     python -m src.cli.curate import-local --sport badminton --video-path "?/Match.MP4" \
79         --annotations-path clean.csv --license Proprietary
80
81     **Do not widen the PK to float** ? the DB column is INTEGER, export orders by it, and
82     the indexer ignores `rally_number` entirely. Resolve ids at the adapter boundary.
83
84     ## Annotator ? canonical field mapping
85
86     | Annotator field | Canonical | Handling |
87     |---|---|---|
88     | `rally_number` (lean) / `rally` (rich) | `rally_number:int` | Lean maps directly;
89     **fractional/dup ids renumbered at the normalizer** (PK is INT) |
90     | `start_time` / `start_mmss` | `start_time:float` | Decimal direct; `mm:ss` converted
91     to seconds |
92     | `end_time` / `end_mmss` | `end_time:float` | **Blank / ? start dropped, not raised** |
93     | `ending_reason` | `ending_reason:str` | Vocab matches the documented canonical set
94     (free-text, not an enforced enum); pipe-delimited folded to primary; **stored but dropped on
95     export** |
96     | `sport` | ? (from `--sport` flag) | CSV column ignored; operator sets the flag |
97     | rich-only `shots`, `true_*`, `VERIFIED`, `notes` | `shots_count` + QA/provenance |
98     Carried where present; none reach the indexer |
99
100     ## Roadmap to the canonical-hub end-state
101
102     All additive, video-level, backward-compatible (consumers ignore unknown manifest keys).
103
104     - **P1 ? persist the labeled window in the manifest.** Partial labels cover only a
105     prefix; a full-video eval penalizes correct detections in the unlabeled tail. The
106     indexer already restricts to a window (`eval_partial_labels.py::restrict_to_window`)
107     but currently *infers* `window_end = max(end_time)` ? which is **wrong when a video is
108     tagged past its last rally** (valid predictions in the labeled-but-rally-free tail get
109     scored as false positives). The collector should own `labeled_window_start/end`
110     (the VLC tool is resumable + monotonic, so the session extent is known), carry it in
111     the manifest, and let the indexer apply it automatically. Short-term: pass
112     `--window-end` at eval time.
113
114     - **P1 ? maturity ladder + provenance.** `CurationStatus` (STAGING/CURATED/REJECTED) is
115     a *workflow* state, not a *trust* state. Add a separate `needs_review` ? confirmed ?
116     gold` enum + `annotation_source` (manual_vlc / manual_review / vendor_cvat) +
117     `annotator_id` + timestamp. Route self-rated imports to `needs_review` instead of the
118     hardcoded `CURATED` (`curate.py`), so the golden-set regression flow restricts the
119     tier to `maturity=gold` and an unreviewed re-tag can't silently become the answer key.

```



```

113     (Today a single-rater self-rating and a vendor-QA'd delivery are indistinguishable.)
114 - **P2 ? unify sources behind one adapter** emitting `(List[RallyRecord], List[Issue])`,
115     and run the **same Issue detectors** (gap>15s possible-missed-rally, shot-density,
116     vocab, overlap) on manual imports that vendor CVAT already gets. Today self-produced
117     labels get *less* scrutiny than a vendor delivery ? backwards.
118 - **P2 ? freeze the intervals-only export contract** in one doc both repos reference +
119     a round-trip test (header matching is exact case-insensitive, so a vendor
120     `start_seconds` would fail silently).
121
122 ## Open question
123
124 Is the **rich** `rally_review.csv` (mm:ss times, a `rally` column with fractional ids,
125 pipe-delimited reasons, `true_` correction columns) ever the *source of truth*, or is
126 the lean VLC `.rallies.csv` always canonical? Neither `import-local` nor
127 `load_canonical_csv` (which keys on `rally_number`) ingest the rich file directly. The
128 `normalize_annotations.py` step already handles its `mm:ss` + `rally` + pipe-reason
129 quirks for the common case; a fuller converter that applies `true_` overrides is only
130 worth building if the rich sheet is canonical.
131

```

8. user (2026-06-10T05:59:57.025Z)

```

1     # Decision records
2
3     Lightweight ADRs for this repo. (The cross-repo training-pipeline ADRs live in
4     the indexer repo's `docs/DECISIONS.md`; this file is for collector-local choices.)
5
6     ---
7
8     ## ADR-001 ? Keep explicit per-sport code instead of a sport registry (for now)
9
10    **Status:** Accepted · **Date:** 2026-06-02
11
12    ### Context
13    Supporting a sport currently requires parallel edits across several sites:
14    - `SEGMENT_TABLES` + a `CREATE TABLE` + `insert_` / `get_` in `src/core/db.py`
15    - a `validate_` method in `src/core/validator.py`
16    - `import-local` parse/validate/insert branches and interactive review/play in
17    `src/cli/curate.py`
18    - a schema in `src/core/schemas/`
19
20    A code review flagged this as copy-paste duplication and proposed a registry/plugin
21    mechanism (e.g. `SPORT_REGISTRY: dict[SportType, SportSpec]` driving generic
22    insert/get/validate/dispatch) to make adding a sport a single entry.
23
24    ### Decision
25    **We are keeping the explicit per-sport code and NOT introducing a registry
26    abstraction at this time.**
27
28    ### Rationale
29    - **It's well-tested and mechanical.** Coverage is ~87%; adding a sport is a known,
30      ~30-minute "mirror an existing sport" task (checklist below), not a research problem.
31    - **Schemas are heterogeneous.** Rally sports (badminton, pickleball, table tennis)
32      share `rally_number` / `shots_count` / `ending_reason`, but tennis (`set_score`,
33      `game_score`, `server`, `ending_type`) and cricket (`over`, `ball`, `runs`,
34      `wicket`, `extras`) differ enough that a single generic table/row mapper would leak
35      special-cases ? trading visible duplication for hidden complexity.
36    - **Explicit code is greppable.** Contributors can follow one sport end-to-end without
37      learning a registry indirection. Premature abstraction is harder to undo than dup.
38    - **Low, slow growth.** Five sports today; new ones arrive rarely.
39
40    ### Consequences
41    - Adding a sport touches ~5 sites (see checklist). We accept that boilerplate.
42    - If the cost rises, ADR-001 should be revisited and superseded.

```

```

43     ### Revisit triggers (when to build the registry)
44     - Sport count grows beyond ~8, **or**
45     - We add a 4th+ rally-identical sport and the dispatch churn starts causing bugs, **or**
46     - A new cross-cutting feature would otherwise require editing every per-sport branch.
47
48     ### Future design sketch (when revisited)
49     ```
50     @dataclass
51     class SportSpec:
52         schema_class: type
53         table_name: str
54         pk_columns: tuple[str, ...]
55         columns: tuple[str, ...]
56         validate_fn: Callable
57
58     SPORT_REGISTRY: dict[SportType, SportSpec] = { ... }
59     # Generic: insert_segments(sport, items) / get_segments(sport, video_id) /
60     #           validate(sport, items) / import-local dispatch driven by the registry.
61     ```
62
63     ---
64
65     ## Checklist: adding a new sport today (under ADR-001)
66
67     1. `src/core/schemas/<sport>.py` ? define the segment model (mirror the closest existing
68     sport) with a `start_time < end_time` validator.
69     2. `src/core/db.py` ? add to `SEGMENT_TABLES`, add a `CREATE TABLE`, `insert_<sport>()`,
70     `get_<sport>()`, and a cascade `DELETE` in `delete_video`.
71     3. `src/core/validator.py` ? add `validate_<sport>()` (duration + no-overlap).
72     4. `src/cli/curate.py` ? add branches in `import_local` (validate + insert),
73     `parse_local_annotations` (JSON **and** CSV), and the interactive review/play.
74     5. `config/config.yaml` ? add the sport to `supported_sports` (and confirm the
75     `SportType` enum value in `schemas/base.py`).
76     6. `tests/` ? mirror the existing db/validator/cli tests for the new sport.

```

9. user (2026-06-10T06:00:02.035Z)

```

1     # Role in the rally-segmenter training pipeline
2
3     This repo is the **data producer** in a two-repo system. Its job is to curate
4     sports videos with per-rally `[start, end)` annotations and commercial-license
5     metadata, then hand that data to the **badminton-highlight-indexer**, which
6     processes, scores, and (Phase 4) **trains** an offline rally segmenter on it.
7
8     This doc explains *why* the `export` and `fetch` commands exist and how they fit
9     the bigger picture. The full decision records and end-to-end narrative live in the
10    indexer repo:
11
12    - `badminton-highlight-indexer/docs/ROADMAP.md` ? the four-phase plan and end-state.
13    - `badminton-highlight-indexer/docs/DECISIONS.md` ? ADR-002 (this bridge), ADR-003
14      (acquisition policy), ADR-004 (offline learned segmenter).
15
16    ## Why this repo matters
17
18    The indexer's best rally mode refines local detections with a **paid Gemini API**.
19    The strategic goal is to replace that with an **offline model trained on our own
20    commercially-licensed, annotated corpus** ? making rally detection free, offline,
21    and self-improving. None of that is possible without a clean, licensable,
22    accurately-labelled, full-resolution dataset. That dataset is what this repo
23    produces.
24
25    ## What the two key commands do (and why)
26
27    - **`curate export`** ? emits the indexer's ground-truth format: per-video

```

```

28     `start,end` CSVs + a `manifest.json` pairing each CSV with its video file.
29     Sport-agnostic (all supported sports ? badminton, tennis, cricket, pickleball,
30     table tennis ? share a `(video_id, ordinal, start, end)` shape) and gated on
31     `is_commercial`, so only licensable data flows
32     downstream. The annotations are temporal-only ? they train/evaluate rally
33     *segmentation*, not the frame-level shuttle detector (see ADR-002).
34
35     - fetch ? (re)downloads videos at best available resolution, updating
36     the DB in place so annotations and curation status are never lost. Uses
37     authenticated cookies for reliable, unthrottled downloads. The why and the
38     hard-won pitfalls (SABR?144p, Chrome DPAPI, rate-limiting) are in ADR-003.
39
40     ## The one contract that matters: video identity
41
42     The indexer keys videos by file MD5, this repo by a human `video_id`. A
43     re-encode/re-download changes the bytes, hence the MD5. So the manifest carries the
44     exact `local_path`, and any downstream processing must run on that file. This is
45     why acquisition (`fetch`) happens before the indexer processes anything.
46

```

10. user (2026-06-10T06:00:02.562Z)

```

1     # Rally Ingestion Pipeline ? End-to-End FAQ
2
3     How a vendor's rally annotations become training/eval ground truth, and how we
4     validate each delivery against our conventions. This is the operational companion
5     to the annotation spec (`docs/annotation/ROORA_BRIEF.md` + `ANNOTATION_GUIDE.md`).
6
7     ```
8     vendor delivery                OUR side (collector)                modeling
(indexer)
9     ?????????????????? convert ?????????????????? validate ??????????????????
10    ? CVAT export ? ?????????????? ? canonical CSV ? ?????????? ? report .md ?
11    ? (boxes + attrs)? ? (per rally, ? ? report .json ?
12    ? OR ? ? decimal secs)? ? review .csv ??? human
verifies
13    ? canonical CSV ? ?????????????????? ??????????????????
flagged rallies
14    ?????????????????? ? ?
15    ? import-local (validated) ?
16    ? golden labels
17    ?????????????????? export ??????????????????
18    ? SQLite (rally ? ?????????? ? <id>.csv ?
start,end (indexer GT)
19    ? tables) ? ? manifest.json?
20    ?????????????????? ??????????????????
21    ?
22    slice clips / calibrate /
train
23    ```
24
25    ## The four commands (in order)
26
27    ```bash
28    # 1. Normalise a CVAT export ? our canonical per-rally CSV (skip if the vendor sent a
CSV)
29    python -m src.cli.curate convert-cvat --input job.xml \
30        --video-path match.mp4 --out canonical.csv --issues issues.txt
31
32    # 2. Validate the delivery ? report + annotatable review sheet (the QA gate)
33    python -m src.cli.curate validate-delivery --input job.xml \
34        --video-path match.mp4 --video-id pilot_badminton_01 --out-dir data/validation
35    # ...or validate a CSV directly:
36    python -m src.cli.curate validate-delivery --csv canonical.csv --video-id
pilot_badminton_01

```

```

37
38 # 3. After human review, import the (corrected) CSV into the DB
39 python -m src.cli.curate import-local --sport badminton \
40     --video-path match.mp4 --annotations-path canonical.csv --fps 59.94
41
42 # 4. Export indexer-ready ground truth (start,end CSVs + manifest)
43 python -m src.cli.curate export --out-dir data/eval_export
44 ```
45
46 ---
47
48 ## FAQ
49
50 ### What delivery formats do we accept?
51 Both. A CVAT export (image- or track-mode XML ? rally fields carried as box
52 attributes) is normalised by `convert-cvat`; an already-canonical CSV is taken
53 as-is. Vendors may keep working in CVAT ? we do not require them to hand-author
54 CSVs.
55
56 ### What is the "canonical" CSV?
57 One row per rally:
58 `rally_number, start_mmss, end_mmss, start_time, end_time, shots_count, ending_reason,
59 last_shot_outcome, score_before, score_after, notes`.
60 `start_time`/`end_time` are decimal seconds; `import-local` reads those.
61
62 ### How are timestamps handled? (the fps thing)
63 We recompute each rally's interval from its frame range × the true fps
64 (`seconds = frame / fps`). A vendor's CVAT time attributes are ignored ? a
65 step-sampled job that assumes 30 fps writes systematically wrong times. fps comes
66 from `--fps` or ffprobe of `--video-path`. `mm:ss` is also accepted everywhere.
67
68 ### What does validation check?
69 | code | severity | meaning |
70 |---|---|---|
71 | `overlap` | blocker | two rallies' intervals overlap (`end[N] > start[N+1]`) |
72 | `zero_duration` | blocker | `end <= start` (single-sample / stray box) |
73 | `frame_out_of_range` | high | a box frame is past the video's last frame |
74 | `shot_density` | medium | long rally with implausibly few shots ? likely an inflated
75 span (often hides a merged rally) |
76 | `possible_missed_rally` | medium | a long unlabeled gap between two rallies ? check
77 for an omitted rally |
78 | `ending_reason_vocab` | medium | `ending_reason` outside the agreed set |
79 | `outcome_vocab` | low | `last_shot_outcome` outside `{in,net,out,n_a}` |
80
81 ### Blocker vs. review ? what's the gate?
82 `validate-delivery` exits non-zero if any blocker exists (so it can gate a bulk
83 loop / CI). Non-blocker issues don't stop the gate but populate the review sheet
84 for the human pass. `import-local` independently rejects overlaps/zero-duration, so
85 bad ground truth can't slip into the DB.
86
87 ### How are missed rallies caught? (the worst error)
88 Two complementary checks, because a miss shows up two ways:
89 1. Clean omission (a real hole in the timeline) ? `possible_missed_rally` (gap
90 heuristic).
91 2. Merged/inflated (the missed rally got swallowed into a neighbour's box, leaving
92 no gap) ? `shot_density` flags the over-long, under-shot rally, and the human finds
93 the second rally inside it.
94 Neither is a guarantee ? they surface regions to review, they don't certify
95 completeness. (A future enhancement: a motion second-opinion over the raw video.)
96
97 ### What's the review sheet?
98 `_review.csv` ? one row per rally with `auto_flags` (which checks fired) plus
99 blank `VERIFIED / true_start_mmss / true_end_mmss / true_ending_reason /
100 true_last_shot_outcome / reviewer_comment` columns. Open it next to the video
101 (timestamp overlay on) and fill the flagged rows. This is the human gate that turns

```

```

98     a delivery into trusted golden labels.
99
100    ### Why two ending fields (`ending_reason` + `last_shot_outcome`)?
101    They're orthogonal: `ending_reason` is *why* the point ended (winner / forced_error /
102    unforced_error / service_fault / let / other); `last_shot_outcome` is *where* the
103    last shot went (in / net / out / n_a). A shuttle that flew long is
104    `forced_error/unforced_error` + `out` ? both kept, no information lost.
105
106    ### How do I onboard a brand-new video?
107    1. Get fps into the intake manifest (we provide it to the vendor).
108    2. `convert-cvat` (if CVAT) ? `validate-delivery` ? read the report, fill the review
109       sheet for flagged rallies ? correct the CSV ? `import-local` ? `export`.
110    3. The same suite runs for every video, so quality stays consistent at scale.
111
112    ### How do I pull a single rally out to eyeball it?
113    The slicer (indexer):
114    ```bash
115    python -m backend.tools.rally_slicer --video match.mp4 \
116        --labels golden_labels.csv --rallies 7,16 --out-dir output/rally_clips
117    ```
118    Clips are padded to match the highlight compiler (0.5s lead / 1.0s tail).
119
120    ### Does coarse frame sampling hurt boundaries?
121    Not in practice: ~0.5s sampling slack is absorbed by the highlight compiler's clip
122    padding, so we don't flag sampling density.
123
124    ### Will a non-commercial license block export?
125    The video is **kept** in the DB regardless of license, but the default `export` is
126    **commercial-only** ? a non-commercial source (e.g. `CC-BY-NC-SA`) is *excluded*
127    from the eval/training set unless you pass `--include-noncommercial`. So if a video
128    "goes missing" downstream, check its `is_commercial` flag first. (The observability
129    report raises a `license_noncommercial` warning for exactly this case.)
130
131    ### What is the observability report? (*.report.json` / `*.report.md` / `*.summary.md`)
132    Every `validate-delivery`, `convert-cvat`, and `import-local` run now also writes a
133    per-video **observability report** (next to the video for `import-local`; into the
134    `--out-dir` / next to the canonical CSV otherwise):
135    - **`<id>.report.json`** ? machine-readable: every stage, metric, and finding. Diff
136      two runs to see what changed between deliveries.
137    - **`<id>.report.md`** ? the **technical** report: parse / validation / registration
138      stages plus **findings** (one per delivery issue, each citing this FAQ). Read this
139      to debug a delivery.
140    - **`<id>.summary.md`** ? a **delivery summary** for the vendor / data team: the
141      rally count and coverage plus the plain-language list of rallies to review ? no
142      internal codes or severities.
143    This is in addition to the existing `<id>_report.md` / `_report.json` / `_review.csv`
144    (the human review sheet), which are unchanged.
145

```

11. user (2026-06-10T06:00:13.288Z)

```

960 C:/Users/avidu/Projects/sports-data-collector/src/cli/curate.py
218 C:/Users/avidu/Projects/sports-data-collector/src/cli/normalize_annotations.py
 49 C:/Users/avidu/Projects/sports-data-collector/src/cli/ingest.py
538 C:/Users/avidu/Projects/sports-data-collector/src/parsers/cvat/rally_cvat.py
567 C:/Users/avidu/Projects/sports-data-collector/src/core/db.py
 49 C:/Users/avidu/Projects/sports-data-collector/src/core/parsing.py
144 C:/Users/avidu/Projects/sports-data-collector/src/core/validator.py
 30 C:/Users/avidu/Projects/sports-data-collector/src/core/config.py
195 C:/Users/avidu/Projects/sports-data-collector/src/core/downloader.py
169 C:/Users/avidu/Projects/sports-data-collector/src/reporting/adapters.py
 21 C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/badminton.py
 40 C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/base.py
 22 C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/cricket.py
 21 C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/pickleball.py

```

```

21 C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/tabletennis.py
20 C:/Users/avidu/Projects/sports-data-collector/src/core/schemas/tennis.py
15 C:/Users/avidu/Projects/sports-data-collector/src/parsers/base_parser.py
150 C:/Users/avidu/Projects/sports-data-collector/src/parsers/badminton/coachai.py
320 C:/Users/avidu/Projects/sports-data-collector/src/crawlers/badminton_crawler.py
301 C:/Users/avidu/Projects/sports-data-collector/tests/test_cli.py
25 C:/Users/avidu/Projects/sports-data-collector/tests/test_config.py
235 C:/Users/avidu/Projects/sports-data-collector/tests/test_cvat_parser.py
247 C:/Users/avidu/Projects/sports-data-collector/tests/test_db.py
145 C:/Users/avidu/Projects/sports-data-collector/tests/test_downloader.py
46 C:/Users/avidu/Projects/sports-data-collector/tests/test_local_import.py
111 C:/Users/avidu/Projects/sports-data-collector/tests/test_normalize_annotations.py
111 C:/Users/avidu/Projects/sports-data-collector/tests/test_parser.py
31 C:/Users/avidu/Projects/sports-data-collector/tests/test_parsing.py
135 C:/Users/avidu/Projects/sports-data-collector/tests/test_reporting.py
49 C:/Users/avidu/Projects/sports-data-collector/tests/test_reporting_faq.py
37 C:/Users/avidu/Projects/sports-data-collector/tests/test_schemas.py
117 C:/Users/avidu/Projects/sports-data-collector/tests/test_validator.py
5139 total

```

12. user (2026-06-10T06:00:18.473Z)

```

1     import os
2     import sys
3     import time
4     import json
5     import csv
6     import argparse
7     from typing import List, Dict, Any, Tuple
8     from tabulate import tabulate
9
10    from src.core.config import load_config
11    from src.core.parsing import safe_int as _safe_int, safe_float_required as _safe_float
12    from src.core.db import SportsDatabase
13    from src.core.schemas.base import VideoMetadata, SportType, LicenseType, CurationStatus
14    from src.core.schemas.badminton import BadmintonRally
15    from src.core.schemas.tennis import TennisPoint
16    from src.core.schemas.cricket import CricketDelivery
17    from src.core.schemas.pickleball import PickleballRally
18    from src.core.schemas.tabletennis import TableTennisRally
19    from src.core.validator import DatasetValidator
20    from src.core.downloader import download_best_video, probe_resolution_fps,
height_to_label
21
22
23    class SportsDataCLI:
24        def __init__(self, config_path: str = "config/config.yaml"):
25            self.config_path = config_path
26            self.config = self._load_config()
27
28            # Resolve DB path
29            db_path = self.config.get("database_path", "data/sports_metadata.db")
30            self.db = SportsDatabase(db_path)
31            self.validator = DatasetValidator(config_path)
32
33        def _load_config(self) -> Dict[str, Any]:
34            if not os.path.exists(self.config_path):
35                print(f"Warning: Configuration file not found at {self.config_path}. Using
defaults.")
36            return load_config(self.config_path)
37
38        def status(self):
39            """Displays status summary of all videos in the database."""
40            # Query total count by status via the public DB read API.
41            rows = self.db.get_status_counts() # list of (sport, status, cnt)

```

```

42
43         if not rows:
44             print("\nDatabase is currently empty. Run crawlers or import local data.\n")
45             return
46
47         data = [[sport, status, cnt] for sport, status, cnt in rows]
48         print("\n=== Dataset Ingestion Status ===")
49         print(tabulate(data, headers=["Sport", "Status", "Count"], tablefmt="grid"))
50         print()
51
52     def import_local(self, args):
53         """Manually registers a local video and annotation file into the database."""
54         sport_str = args.sport.lower()
55         if sport_str not in self.config.get("supported_sports", ["badminton"]):
56             print(f"Error: Sport '{sport_str}' is not supported. Supported:
57 {self.config.get('supported_sports')}")
58             sys.exit(1)
59
60         # Validate video path
61         if not os.path.exists(args.video_path):
62             print(f"Warning: Local video path '{args.video_path}' does not exist on
63 disk, but continuing registration.")
64
65         # Validate annotation file
66         if not os.path.exists(args.annotations_path):
67             print(f"Error: Annotations file '{args.annotations_path}' not found.")
68             sys.exit(1)
69
70         match_name = args.match_name or os.path.basename(args.video_path).split('.')[0]
71         video_id = args.video_id or match_name.lower().replace(" ", "_")
72
73         # Parse license
74         try:
75             license_type = LicenseType(args.license)
76         except ValueError:
77             print(f"Warning: Unknown license '{args.license}'. Defaulting to Unknown
78 (non-commercial).")
79             license_type = LicenseType.UNKNOWN
80
81         is_commercial = self.validator.check_license_commercial(license_type)
82
83         # 1. Parse Annotations
84         annotations = []
85         try:
86             annotations = self._parse_local_annotations(sport_str, video_id,
87 args.annotations_path)
88         except Exception as e:
89             print(f"Error parsing annotations: {e}")
90             sys.exit(1)
91
92         # 2. Run Validation
93         is_valid = True
94         errors = []
95         if sport_str == "badminton":
96             is_valid, errors = self.validator.validate_badminton_rallies(annotations)
97         elif sport_str == "tennis":
98             is_valid, errors = self.validator.validate_tennis_points(annotations)
99         elif sport_str == "cricket":
100             is_valid, errors = self.validator.validate_cricket_deliveries(annotations)
101         elif sport_str == "pickleball":
102             is_valid, errors = self.validator.validate_pickleball_rallies(annotations)
103         elif sport_str == "tabletennis":
104             is_valid, errors = self.validator.validate_tabletennis_rallies(annotations)
105
106         if not is_valid:

```

```

103         print(f"Error: Annotations failed validation!")
104     for err in errors:
105         print(f" - {err}")
106     sys.exit(1)
107
108 # 3. Create Video Metadata
109 video_meta = VideoMetadata(
110     video_id=video_id,
111     sport=SportType(sport_str),
112     video_url=None,
113     local_path=args.video_path,
114     license_type=license_type,
115     license_url=None,
116     is_commercial=is_commercial,
117     status=CurationStatus.CURATED, # Custom local imports bypass staging by
default
118     match_name=match_name,
119     fps=args.fps,
120     resolution=args.resolution
121 )
122
123 # 4. Insert into Database
124 self.db.insert_video(video_meta)
125
126 count = 0
127 if sport_str == "badminton":
128     count = self.db.insert_badminton_rallies(annotations)
129 elif sport_str == "tennis":
130     count = self.db.insert_tennis_points(annotations)
131 elif sport_str == "cricket":
132     count = self.db.insert_cricket_deliveries(annotations)
133 elif sport_str == "pickleball":
134     count = self.db.insert_pickleball_rallies(annotations)
135 elif sport_str == "tabletennis":
136     count = self.db.insert_tabletennis_rallies(annotations)
137
138 print(f"\nSuccessfully imported local video '{match_name}':")
139 print(f" - Video ID: {video_id}")
140 print(f" - Sport: {sport_str}")
141 print(f" - Commercial Use Allowed: {is_commercial} (License:
{license_type.value})")
142 print(f" - Registered {count} annotated segments.\n")
143
144 # Per-video observability report (next to the video). Never raises.
145 from src.reporting import emit_import_report
146 src_fmt = os.path.splitext(args.annotations_path)[1].lstrip(".").lower() or
"unknown"
147 report_paths = emit_import_report(
148     video_meta_obj=video_meta, annotations=annotations, is_valid=True,
149     errors=[], count=count, sport=sport_str, source_format=src_fmt,
150     fallback_dir=os.path.dirname(os.path.abspath(args.annotations_path)),
151 )
152 if report_paths:
153     print(f" - Observability report: {report_paths.get('technical_md')}")
154     print(f" - Customer summary      : {report_paths.get('customer_md')}\n")
155
156 def _parse_local_annotations(self, sport: str, video_id: str, filepath: str) ->
List[Any]:
157     """Parses simple JSON or CSV annotations files."""
158     parsed_items = []
159     ext = os.path.splitext(filepath)[1].lower()
160
161     if ext == '.json':
162         with open(filepath, 'r', encoding='utf-8') as f:
163             data = json.load(f)

```



```

164
165         if not isinstance(data, list):
166             raise ValueError("JSON file must contain a list of objects.")
167
168     # Uses the same safe casts as the CSV path: a missing/blank required
169     # time raises a clear ValueError (caught by the caller) instead of a
170     # bare KeyError, and optional ints degrade to sensible defaults.
171     for idx, obj in enumerate(data):
172         if sport == "badminton":
173             parsed_items.append(BadmintonRally(
174                 video_id=video_id,
175                 rally_number=_safe_int(obj.get("rally_number"), idx + 1),
176                 start_time=_safe_float(obj.get("start_time")),
177                 end_time=_safe_float(obj.get("end_time")),
178                 score_before=obj.get("score_before"),
179                 score_after=obj.get("score_after"),
180                 shots_count=_safe_int(obj.get("shots_count"), None),
181                 ending_reason=obj.get("ending_reason"),
182                 last_shot_outcome=obj.get("last_shot_outcome")
183             ))
184         elif sport == "tennis":
185             parsed_items.append(TennisPoint(
186                 video_id=video_id,
187                 point_number=_safe_int(obj.get("point_number"), idx + 1),
188                 start_time=_safe_float(obj.get("start_time")),
189                 end_time=_safe_float(obj.get("end_time")),
190                 set_score=obj.get("set_score"),
191                 game_score=obj.get("game_score"),
192                 server=obj.get("server"),
193                 ending_type=obj.get("ending_type")
194             ))
195         elif sport == "cricket":
196             parsed_items.append(CricketDelivery(
197                 video_id=video_id,
198                 over=_safe_int(obj.get("over"), 0),
199                 ball=_safe_int(obj.get("ball"), 0),
200                 start_time=_safe_float(obj.get("start_time")),
201                 end_time=_safe_float(obj.get("end_time")),
202                 batsman=obj.get("batsman"),
203                 bowler=obj.get("bowler"),
204                 runs=_safe_int(obj.get("runs"), None),
205                 wicket=bool(obj.get("wicket", False)),
206                 extras=_safe_int(obj.get("extras"), 0)
207             ))
208         elif sport == "pickleball":
209             parsed_items.append(PickleballRally(
210                 video_id=video_id,
211                 rally_number=_safe_int(obj.get("rally_number"), idx + 1),
212                 start_time=_safe_float(obj.get("start_time")),
213                 end_time=_safe_float(obj.get("end_time")),
214                 score_before=obj.get("score_before"),
215                 score_after=obj.get("score_after"),
216                 shots_count=_safe_int(obj.get("shots_count"), None),
217                 ending_reason=obj.get("ending_reason"),
218                 last_shot_outcome=obj.get("last_shot_outcome")
219             ))
220         elif sport == "tabletennis":
221             parsed_items.append(TableTennisRally(
222                 video_id=video_id,
223                 rally_number=_safe_int(obj.get("rally_number"), idx + 1),
224                 start_time=_safe_float(obj.get("start_time")),
225                 end_time=_safe_float(obj.get("end_time")),
226                 score_before=obj.get("score_before"),
227                 score_after=obj.get("score_after"),
228                 shots_count=_safe_int(obj.get("shots_count"), None),

```

```

229                 ending_reason=obj.get("ending_reason"),
230                 last_shot_outcome=obj.get("last_shot_outcome")
231             ))
232     elif ext == '.csv':
233         with open(filepath, 'r', encoding='utf-8') as f:
234             reader = csv.DictReader(f)
235             # Normalize headers
236             reader.fieldnames = [n.lower().strip() for n in reader.fieldnames] if
reader.fieldnames else []
237
238         for idx, row in enumerate(reader):
239             if sport == "badminton":
240                 parsed_items.append(BadmintonRally(
241                     video_id=video_id,
242                     rally_number=_safe_int(row.get("rally_number"), idx + 1),
243                     start_time=_safe_float(row.get("start_time")),
244                     end_time=_safe_float(row.get("end_time")),
245                     score_before=row.get("score_before"),
246                     score_after=row.get("score_after"),
247                     shots_count=_safe_int(row.get("shots_count"), None),
248                     ending_reason=row.get("ending_reason"),
249                     last_shot_outcome=row.get("last_shot_outcome")
250                 ))
251             elif sport == "tennis":
252                 parsed_items.append(TennisPoint(
253                     video_id=video_id,
254                     point_number=_safe_int(row.get("point_number"), idx + 1),
255                     start_time=_safe_float(row.get("start_time")),
256                     end_time=_safe_float(row.get("end_time")),
257                     set_score=row.get("set_score"),
258                     game_score=row.get("game_score"),
259                     server=row.get("server"),
260                     ending_type=row.get("ending_type")
261                 ))
262             elif sport == "cricket":
263                 parsed_items.append(CricketDelivery(
264                     video_id=video_id,
265                     over=_safe_int(row.get("over"), 0),
266                     ball=_safe_int(row.get("ball"), 0),
267                     start_time=_safe_float(row.get("start_time")),
268                     end_time=_safe_float(row.get("end_time")),
269                     batsman=row.get("batsman"),
270                     bowler=row.get("bowler"),
271                     runs=_safe_int(row.get("runs"), None),
272                     wicket=(row.get("wicket") or "").lower() in ["true", "1",
"yes"],
273                     extras=_safe_int(row.get("extras"), 0)
274                 ))
275             elif sport == "pickleball":
276                 parsed_items.append(PickleballRally(
277                     video_id=video_id,
278                     rally_number=_safe_int(row.get("rally_number"), idx + 1),
279                     start_time=_safe_float(row.get("start_time")),
280                     end_time=_safe_float(row.get("end_time")),
281                     score_before=row.get("score_before"),
282                     score_after=row.get("score_after"),
283                     shots_count=_safe_int(row.get("shots_count"), None),
284                     ending_reason=row.get("ending_reason"),
285                     last_shot_outcome=row.get("last_shot_outcome")
286                 ))
287             elif sport == "tabletennis":
288                 parsed_items.append(TableTennisRally(
289                     video_id=video_id,
290                     rally_number=_safe_int(row.get("rally_number"), idx + 1),
291                     start_time=_safe_float(row.get("start_time")),

```

```

292             end_time=_safe_float(row.get("end_time")),
293             score_before=row.get("score_before"),
294             score_after=row.get("score_after"),
295             shots_count=_safe_int(row.get("shots_count"), None),
296             ending_reason=row.get("ending_reason"),
297             last_shot_outcome=row.get("last_shot_outcome")
298         ))
299     else:
300         raise ValueError("Unsupported annotations file extension. Use .json or
.csv.")
301
302     return parsed_items
303
304 def _fetch_one_with_retry(self, url, dest_path, args, max_height, max_size_mb):
305     """Download one video at best resolution, retrying transient failures.
306
307     A below-min-height result almost always means YouTube rate-limited the
308     request and served a low fallback (not that the video lacks HD), so we
309     retry with a growing backoff. Returns (DownloadResult, None) on success,
310     or (None, failure_reason) once retries are exhausted. A below-min file is
311     discarded each attempt so it never lands in the DB.
312     """
313     attempts = max(1, args.retries + 1)
314     reason = "unknown error"
315     for i in range(attempts):
316         result = download_best_video(
317             url, dest_path,
318             max_height=max_height, max_size_mb=max_size_mb,
319             cookies_from_browser=args.cookies_from_browser,
320             cookies_file=args.cookies,
321             player_client=args.player_client,
322         )
323         if result.success and (result.height is None or result.height >=
args.min_height):
324             return result, None
325
326         if not result.success:
327             reason = result.error or "unknown error"
328         else:
329             reason = (f"got {result.resolution_label} (< min-height
{args.min_height})p) ? "
330                     f"YouTube served a low format (rate-limited?)")
331         try:
332             os.remove(result.path)
333         except OSError:
334             pass
335
336         if i < attempts - 1:
337             backoff = args.retry_backoff * (i + 1)
338             print(f" attempt {i+1}/{attempts} failed: {reason}\n retrying in
{backoff:.0f}s...")
339             time.sleep(backoff)
340     return None, reason
341
342 def fetch(self, args):
343     """(Re)download videos at best available resolution, updating the DB in place.
344
345     This is the in-place upgrade path: rows, rally annotations, and curation
346     status are all preserved ? only the video file plus its `local_path`,
347     `resolution`, and `fps` are refreshed. By default it skips videos that
348     already have a local file at or above `--min-height`, so re-runs are cheap
349     and resumable; `--force` re-fetches everything.
350     """
351     storage = self.config.get("video_storage", {})
352     local_dir = storage.get("local_dir", "./data/videos")

```

```

353         max_download_size_mb = storage.get("max_download_size_mb")
354         max_height = args.max_height if args.max_height is not None else
storage.get("max_height")
355         os.makedirs(local_dir, exist_ok=True)
356
357         commercial_only = not args.include_noncommercial
358         if args.status == "all":
359             statuses = [CurationStatus.CURATED, CurationStatus.STAGING]
360         else:
361             statuses = [CurationStatus(args.status)]
362
363         supported = self.config.get("supported_sports", ["badminton"])
364         sport_strs = [args.sport] if args.sport else supported
365
366         # Collect target videos.
367         videos = []
368         for sport_str in sport_strs:
369             try:
370                 sport = SportType(sport_str.lower())
371             except ValueError:
372                 print(f"Warning: skipping unknown sport '{sport_str}'.")
373                 continue
374             for status in statuses:
375                 videos.extend(self.db.get_videos_by_status(status, sport))
376         if args.video_id:
377             videos = [v for v in videos if v.video_id == args.video_id]
378
379         if not videos:
380             print("\nNo matching videos to fetch.\n")
381             return
382
383         results = [] # (video_id, action, detail)
384         for video in videos:
385             if commercial_only and not video.is_commercial:
386                 results.append((video.video_id, "skip", "non-commercial"))
387                 continue
388             if not video.video_url or "watch?v=example" in video.video_url:
389                 # Either no URL, or the parser's placeholder (the ShuttleSet catalog
390                 # had no URL for this match) ? nothing to download.
391                 results.append((video.video_id, "skip", "no real source URL"))
392                 continue
393
394             dest_path = os.path.join(local_dir, f"{video.video_id}.mp4")
395
396             # Skip if we already have a good-enough local file (unless --force).
397             if not args.force and os.path.exists(dest_path):
398                 _, height, _ = probe_resolution_fps(dest_path)
399                 if height is not None and height >= args.min_height:
400                     results.append((video.video_id, "have", f"{height_to_label(height)}
already"))
401                     continue
402
403             print(f"Fetching {video.video_id} ({video.match_name})...")
404             result, fail_reason = self._fetch_one_with_retry(
405                 video.video_url, dest_path, args, max_height, max_download_size_mb)
406             if fail_reason is not None:
407                 results.append((video.video_id, "FAIL", fail_reason))
408                 if args.sleep:
409                     time.sleep(args.sleep)
410                 continue
411
412             # Update only the file-derived fields; preserve everything else (UPSERT
413             # mutates the row in place, leaving rallies intact).
414             video.local_path = result.path
415             if result.resolution_label:

```

```

416         video.resolution = result.resolution_label
417     if result.fps:
418         video.fps = result.fps
419     self.db.insert_video(video)
420
421     detail = (f"{result.resolution_label or '?'}"
422              f"{f' {result.fps:g}fps' if result.fps else ''},
{result.size_mb:.0f} MB")
423     results.append((video.video_id, "fetched", detail))
424     if args.sleep:
425         time.sleep(args.sleep)
426
427     print("\n=== Fetch summary ===")
428     print(tabulate([[vid, action, detail] for vid, action, detail in results],
429                   headers=["Video ID", "Action", "Detail"], tablefmt="grid"))
430     fetched = sum(1 for _, a, _ in results if a == "fetched")
431     failed = [r for r in results if r[1] == "FAIL"]
432     print(f"\nFetched/updated : {fetched}    Already-good : {sum(1 for _,a,_ in
results if a=='have')}}    "
433          f"Skipped : {sum(1 for _,a,_ in results if a=='skip')}}    Failed :
{len(failed)}}")
434     if failed:
435         print("Failures (left unchanged in DB):")
436         for vid, _, detail in failed:
437             print(f" - {vid}: {detail}")
438     print()
439
440 def export_eval_set(self, args):
441     """Export curated annotations as indexer-ready eval/training sets.
442
443     Emits, per qualifying video, a `

```

```

478         if commercial_only and not video.is_commercial:
479             skipped_noncomm += 1
480             continue
481
482         intervals = self.db.get_segment_intervals(video.video_id, sport)
483         if not intervals:
484             skipped_empty += 1
485             continue
486
487         sport_dir = os.path.join(out_dir, sport.value)
488         os.makedirs(sport_dir, exist_ok=True)
489         csv_path = os.path.join(sport_dir, f"{video.video_id}.csv")
490         with open(csv_path, "w", newline="", encoding="utf-8") as f:
491             writer = csv.writer(f)
492             writer.writerow(["start", "end"])
493             for start, end in intervals:
494                 writer.writerow([start, end])
495
496         manifest.append({
497             "video_id": video.video_id,
498             "sport": sport.value,
499             "csv": os.path.relpath(csv_path, out_dir).replace(os.sep, "/"),
500             "local_path": video.local_path,
501             "video_url": video.video_url,
502             "match_name": video.match_name,
503             "fps": video.fps,
504             "resolution": video.resolution,
505             "license_type": video.license_type.value,
506             "is_commercial": video.is_commercial,
507             "status": video.status.value,
508             "num_segments": len(intervals),
509         })
510
511     os.makedirs(out_dir, exist_ok=True)
512     manifest_doc = {
513         "commercial_only": commercial_only,
514         "statuses": [s.value for s in statuses],
515         "total_videos": len(manifest),
516         "total_segments": sum(m["num_segments"] for m in manifest),
517         "eval_sets": manifest,
518     }
519     manifest_path = os.path.join(out_dir, "manifest.json")
520     with open(manifest_path, "w", encoding="utf-8") as f:
521         json.dump(manifest_doc, f, indent=2)
522
523     print(f"\n=== Exported eval/training set to '{out_dir}' ===")
524     if manifest:
525         summary = [[m["sport"], m["video_id"], m["num_segments"],
526                     "YES" if m["is_commercial"] else "NO",
527                     os.path.basename(m["local_path"]) if m["local_path"] else "(no
local file)"]
528                    for m in manifest]
529         print(tabulate(summary, headers=["Sport", "Video ID", "Segments", "Comm?",
"Video file"], tablefmt="grid"))
530         print(f"\nVideos exported : {len(manifest)}")
531         print(f"Total segments : {manifest_doc['total_segments']}")
532         if skipped_noncomm:
533             print(f"Skipped (non-commercial) : {skipped_noncomm} (pass --include-
noncommercial to include)")
534         if skipped_empty:
535             print(f"Skipped (no annotations) : {skipped_empty}")
536         print(f"Manifest : {manifest_path}\n")
537
538     def convert_cvat(self, args):
539         """Convert a vendor CVAT rally export into our canonical per-rally CSV.

```

```

540
541 Absorbs the CVAT bounding-box delivery format (rally fields carried as box
542 attributes) on the ingestion side: each rally's time is recomputed from
543 its frame range at the video's true fps, and content defects (overlaps,
544 zero-duration, off-vocabulary endings, coarse sampling) are written to an
545 issues report rather than silently "fixed". The emitted CSV imports via
546 `import-local`; any blocker issue must be resolved first, since the
547 importer's validator rejects overlapping / zero-duration rallies.
548 """
549 from src.parsers.cvat.rally_cvat import (
550     convert, write_canonical_csv, render_issue_report,
551 )
552 try:
553     result = convert(args.input, fps=args.fps,
554                      video_path=args.video_path, step=args.step)
555 except (FileNotFoundError, ValueError) as e:
556     print(f"Error: {e}")
557     sys.exit(1)
558
559 write_canonical_csv(result.records, args.out)
560
561 # Observability report next to the canonical CSV. Never raises.
562 from src.reporting import emit_delivery_report
563 src_name = os.path.splitext(os.path.basename(args.input))[0]
564 dur = (result.meta.get("stop_frame") / result.fps) if
565 (result.meta.get("stop_frame") and result.fps) else None
566 emit_delivery_report(
567     out_dir=os.path.dirname(os.path.abspath(args.out)) or ".", source=src_name,
568     records=result.records, issues=result.issues, fps=result.fps,
569     video_path=args.video_path,
570     video_duration_s=dur, source_format="cvat", step=result.step,
571 )
572
573 report = render_issue_report(result, os.path.basename(args.input))
574 if args.issues:
575     with open(args.issues, "w", encoding="utf-8") as f:
576         f.write(report)
577     print(report)
578     print(f"Canonical CSV written: {args.out} ({len(result.records)} rallies)")
579 if result.n_blockers:
580     print(f"\n {result.n_blockers} BLOCKER issue(s) ? import-local will reject
581 "
582         f"this file until they are resolved (by the vendor or manual
583 review).")
584 else:
585     print("\n No blockers. Next:")
586     print(f"    python -m src.cli.curate import-local --sport <sport> "
587           f"--video-path <video> --annotations-path {args.out} --fps
588 {result.fps:g}")
589
590 def validate_delivery(self, args):
591     """Run the full validation suite on a rally delivery and emit a report.
592
593     Accepts a CVAT export (--input + --video-path/--fps) or an already-canonical
594     CSV (--csv). Writes, into --out-dir:
595     <id>_report.md    human report (scoreboard + issues + rallies to review)
596     <id>_report.json  machine summary (for CI / a bulk-ingest loop)
597     <id>_review.csv   the annotatable manual-verification sheet
598     Exits non-zero if any blocker issue exists, so it can gate a bulk pipeline.
599     """
600     import json
601     from src.parsers.cvat.rally_cvat import (
602         convert, load_canonical_csv, detect_issues,
603         build_summary, render_validation_md, write_review_csv,
604     )

```

```

600     vid = args.video_id or "delivery"
601     fps = None
602     step = None
603     stop_frame = None
604     if args.input:
605         try:
606             result = convert(args.input, fps=args.fps, video_path=args.video_path)
607         except (FileNotFoundError, ValueError) as e:
608             print(f"Error: {e}")
609             sys.exit(1)
610         records, issues, fps = result.records, result.issues, result.fps
611         step = result.step
612         stop_frame = result.meta.get("stop_frame")
613     elif args.csv:
614         try:
615             records = load_canonical_csv(args.csv)
616         except FileNotFoundError as e:
617             print(f"Error: {e}")
618             sys.exit(1)
619         if not records:
620             print(f"Error: no usable rally rows in {args.csv}")
621             sys.exit(1)
622         issues = detect_issues(records)
623     else:
624         print("Error: provide --input <cvat.xml> or --csv <canonical.csv>")
625         sys.exit(1)
626
627     os.makedirs(args.out_dir, exist_ok=True)
628     summary = build_summary(records, issues, source=vid, fps=fps)
629     md = render_validation_md(records, issues, source=vid)
630     json_path = os.path.join(args.out_dir, f"{vid}_report.json")
631     md_path = os.path.join(args.out_dir, f"{vid}_report.md")
632     review_path = os.path.join(args.out_dir, f"{vid}_review.csv")
633     with open(json_path, "w", encoding="utf-8") as f:
634         json.dump(summary, f, indent=2)
635     with open(md_path, "w", encoding="utf-8") as f:
636         f.write(md)
637     write_review_csv(records, issues, review_path)
638
639     # Observability report (envelope JSON + technical md + customer summary),
640     # colocated with the existing artifacts. Never raises.
641     from src.reporting import emit_delivery_report
642     duration_s = (stop_frame / fps) if (stop_frame and fps) else None
643     report_paths = emit_delivery_report(
644         out_dir=args.out_dir, source=vid, records=records, issues=issues, fps=fps,
645         video_path=args.video_path, video_duration_s=duration_s,
646         source_format=("cvat" if args.input else "csv"), step=step,
647     )
648
649     print(md)
650     print(f"\nWrote:\n {md_path}\n {json_path}\n {review_path}")
651     if report_paths:
652         print(f" {report_paths.get('technical_md')}\n
653 {report_paths.get('customer_md')}\n {report_paths.get('json')}")
654         if summary["n_blockers"]:
655             print(f"\n GATE: BLOCKED ? {summary['n_blockers']} blocker(s) must be
656 resolved.")
657             sys.exit(2)
658         print("\n GATE: PASS ? review the flagged rallies, then `import-local`.")
659
660     def curate(self):
661         """Interactive loop to review and curate staging video annotations."""
662         staging_videos = self.db.get_videos_by_status(CurationStatus.STAGING)
663         if not staging_videos:
664             print("\nNo staging video annotations to review.\n")

```



```

663         return
664
665     print(f"\nFound {len(staging_videos)} matches in STAGING state for curation:")
666     table_data = []
667     for idx, video in enumerate(staging_videos):
668         table_data.append([
669             idx + 1,
670             video.video_id,
671             video.match_name,
672             video.sport.value,
673             video.license_type.value,
674             "YES" if video.is_commercial else "NO"
675         ])
676
677     print(tabulate(table_data, headers=["#", "Video ID", "Match Name", "Sport",
678 "License", "Commercial Safe?"], tablefmt="grid"))
679     print("\nEnter a number (1-{}) to review details, or 'q' to
680 quit:".format(len(staging_videos)))
681
682     choice = input("> ").strip()
683     if choice.lower() == 'q':
684         return
685
686     try:
687         idx = int(choice) - 1
688         if idx < 0 or idx >= len(staging_videos):
689             print("Invalid choice.")
690             return
691         self._curate_video_interactive(staging_videos[idx])
692     except ValueError:
693         print("Invalid input.")
694
695 def _curate_video_interactive(self, video: VideoMetadata):
696     """Displays details of a specific staging video and prompts for curation
697     action."""
698
699     # Load and count events
700     rallies = []
701     pts = []
702     delivs = []
703     if video.sport == SportType.BADMINTON:
704         rallies = self.db.get_badminton_rallies(video.video_id)
705     elif video.sport == SportType.PICKLEBALL:
706         rallies = self.db.get_pickleball_rallies(video.video_id)
707     elif video.sport == SportType.TABLE_TENNIS:
708         rallies = self.db.get_tabletennis_rallies(video.video_id)
709     elif video.sport == SportType.TENNIS:
710         pts = self.db.get_tennis_points(video.video_id)
711     elif video.sport == SportType.CRICKET:
712         delivs = self.db.get_cricket_deliveries(video.video_id)
713
714     while True:
715         print("\n===== REVIEWING VIDEO =====")
716         print(f"Match Name      : {video.match_name}")
717         print(f"Video ID       : {video.video_id}")
718         print(f"Sport          : {video.sport.value}")
719         print(f"Video URL      : {video.video_url or 'N/A'}")
720         print(f"Local Path     : {video.local_path or 'N/A'}")
721         print(f"License        : {video.license_type.value} ({video.license_url or 'no
722 URL'})")
723
724         print(f"Commercial OK: {video.is_commercial}")
725
726         if video.sport in (SportType.BADMINTON, SportType.PICKLEBALL,
727 SportType.TABLE_TENNIS):
728             print(f"Total Rallies: {len(rallies)}")

```

```

723         if rallies:
724             print("Sample rallies:")
725             sample_data = [[r.rally_number, f"{r.start_time}s - {r.end_time}s",
r.score_after, r.shots_count] for r in rallies[:5]]
726             print(tabulate(sample_data, headers=["Rally #", "Duration", "Score",
"Shots"], tablefmt="simple"))
727         elif video.sport == SportType.TENNIS:
728             print(f"Total Points : {len(pts)}")
729             if pts:
730                 print("Sample points:")
731                 sample_data = [[p.point_number, f"{p.start_time}s - {p.end_time}s",
p.game_score] for p in pts[:5]]
732                 print(tabulate(sample_data, headers=["Point #", "Duration",
"Score"], tablefmt="simple"))
733             elif video.sport == SportType.CRICKET:
734                 print(f"Total Balls : {len(delivs)}")
735                 if delivs:
736                     print("Sample deliveries:")
737                     sample_data = [[f"{d.over}.{d.ball}", f"{d.start_time}s -
{d.end_time}s", d.runs, "Wicket" if d.wicket else ""] for d in delivs[:5]]
738                     print(tabulate(sample_data, headers=["Over.Ball", "Duration",
"Runs", "Wicket"], tablefmt="simple"))
739
740             print("\nChoose Curation Action:")
741             print(" [a] Approve (promote to CURATED)")
742             print(" [r] Reject (delete from database)")
743             print(" [p] Play a rally/point segment in browser")
744             print(" [k] Keep in Staging (do nothing and go back)")
745
746             action = input("> ").strip().lower()
747             if action == 'a':
748                 self.db.update_status(video.video_id, CurationStatus.CURATED)
749                 print(f"Successfully promoted '{video.match_name}' to CURATED.\n")
750                 break
751             elif action == 'r':
752                 self.db.delete_video(video.video_id)
753                 print(f"Successfully rejected and deleted '{video.match_name}'.\n")
754                 break
755             elif action == 'k':
756                 print("Curation state unchanged.\n")
757                 break
758             elif action == 'p':
759                 self._play_segment_interactive(video, rallies, pts, delivs)
760             else:
761                 print("Invalid action selected.")
762
763     def _play_segment_interactive(self, video: VideoMetadata, rallies:
List[BadmintonRally], pts: List[TennisPoint], delivs: List[CricketDelivery]):
764         import webbrowser
765
766         # Get start/end based on user input
767         start_time = None
768         end_time = None
769         label = "segment"
770
771         if video.sport in (SportType.BADMINTON, SportType.PICKLEBALL,
SportType.TABLE_TENNIS):
772             if not rallies:
773                 print("No rallies to play.")
774                 return
775             idx_str = input(f"Enter Rally number to play (1-{len(rallies)}): ").strip()
776             try:
777                 rally_num = int(idx_str)
778                 rally = next((r for r in rallies if r.rally_number == rally_num), None)
779                 if rally:

```

```

780         start_time = rally.start_time
781         end_time = rally.end_time
782         label = f"Rally {rally_num}"
783     except ValueError:
784         pass
785 elif video.sport == SportType.TENNIS:
786     if not pts:
787         print("No points to play.")
788         return
789     idx_str = input(f"Enter Point number to play (1-{len(pts)}): ").strip()
790     try:
791         pt_num = int(idx_str)
792         pt = next((p for p in pts if p.point_number == pt_num), None)
793         if pt:
794             start_time = pt.start_time
795             end_time = pt.end_time
796             label = f"Point {pt_num}"
797     except ValueError:
798         pass
799 elif video.sport == SportType.CRICKET:
800     if not delivs:
801         print("No deliveries to play.")
802         return
803     over_str = input("Enter Over number: ").strip()
804     ball_str = input("Enter Ball number: ").strip()
805     try:
806         o_num = int(over_str)
807         b_num = int(ball_str)
808         d = next((x for x in delivs if x.over == o_num and x.ball == b_num),
809 None)
810         if d:
811             start_time = d.start_time
812             end_time = d.end_time
813             label = f"Over {o_num} Ball {b_num}"
814     except ValueError:
815         pass
816 if start_time is None or end_time is None:
817     print("Segment not found or invalid input.")
818     return
819
820 print(f"Playing {label} (start: {start_time}s, end: {end_time}s)...")
821
822 # Check video location
823 play_url = None
824 if video.local_path and os.path.exists(video.local_path):
825     abs_path = os.path.abspath(video.local_path)
826     # Create file:// URL with media fragment (browsers support #t=start,end
827 fragment)
828     play_url = f"file:/// {abs_path.replace(os.sep,
829 '/' )}#t={start_time},{end_time}"
830 elif video.video_url:
831     # YouTube URL: seek to start time
832     if "youtube.com" in video.video_url or "youtu.be" in video.video_url:
833         t_sec = int(start_time)
834         sep = "&" if "?" in video.video_url else "?"
835         play_url = f"{video.video_url}{sep}t={t_sec}s"
836     else:
837         play_url = video.video_url
838
839 if play_url:
840     print(f"Opening browser: {play_url}")
841     webbrowser.open(play_url)
842 else:
843     print("Error: No local video file or remote URL available to play.")

```

```

842
843
844     def main():
845         parser = argparse.ArgumentParser(description="Sports Data Curator CLI")
846         subparsers = parser.add_subparsers(dest="command")
847
848         # Status command
849         subparsers.add_parser("status", help="Display summary status of sports metadata
database")
850
851         # Curate command
852         subparsers.add_parser("curate", help="Launch interactive curation tool for staging
entries")
853
854         # Import-local command
855         import_parser = subparsers.add_parser("import-local", help="Manually import a local
video and annotations")
856         import_parser.add_argument("--sport", required=True, help="Sport type (e.g.
badminton, tennis, cricket)")
857         import_parser.add_argument("--video-path", required=True, help="Path to local video
file")
858         import_parser.add_argument("--annotations-path", required=True, help="Path to local
JSON/CSV annotations file")
859         import_parser.add_argument("--match-name", help="Name of the match (defaults to
video filename)")
860         import_parser.add_argument("--video-id", help="Custom alphanumeric video ID")
861         import_parser.add_argument("--license", default="Proprietary", help="License type
(e.g. MIT, CC-BY-NC, Proprietary)")
862         import_parser.add_argument("--fps", type=float, default=30.0, help="Frames per
second of the video")
863         import_parser.add_argument("--resolution", default="1080p", help="Resolution of the
video (e.g., 1080p, 720p)")
864
865         # Fetch command ? (re)download videos at best resolution, update DB in place
866         fetch_parser = subparsers.add_parser(
867             "fetch", help="(Re)download videos at best available resolution, updating the DB
in place")
868         fetch_parser.add_argument("--sport", help="Restrict to a single sport (default: all
supported_sports)")
869         fetch_parser.add_argument("--video-id", help="Fetch a single video by ID")
870         fetch_parser.add_argument("--status", choices=["staging", "curated", "all"],
default="all",
871                                     help="Which curation statuses to fetch (default: all)")
872         fetch_parser.add_argument("--min-height", type=int, default=720,
873                                     help="Skip videos already at/above this height unless
--force (default: 720)")
874         fetch_parser.add_argument("--max-height", type=int, default=None,
875                                     help="Cap download resolution (e.g. 1080); default: best
available")
876         fetch_parser.add_argument("--force", action="store_true",
877                                     help="Re-download even if a good-enough local file already
exists")
878         fetch_parser.add_argument("--include-noncommercial", action="store_true",
879                                     help="Include videos whose license is not commercially
safe (default: commercial only)")
880         fetch_parser.add_argument("--cookies-from-browser", dest="cookies_from_browser",
metavar="BROWSER",
881                                     help="Use a logged-in browser's YouTube cookies (e.g.
chrome, edge, firefox, brave). "
882                                     "A Premium login gives reliable best-resolution
downloads. "
883                                     "Chromium browsers must be CLOSED on Windows (cookie
DB lock).")
884         fetch_parser.add_argument("--cookies", metavar="FILE",
885                                     help="Path to a Netscape-format cookies.txt (alternative

```

```

to --cookies-from-browser)")
886     fetch_parser.add_argument("--player-client", dest="player_client",
default="default", metavar="CLIENT",
887                                     help="YouTube player client yt-dlp uses (default:
'default'). Avoid web/web_safari "
888                                     "(SABR-forced -> drops to 144p). 'tv'/'mweb' also
serve 1080p if needed.")
889     fetch_parser.add_argument("--sleep", type=float, default=4.0, metavar="SECONDS",
890                                     help="Pause between videos to avoid YouTube rate-limiting
(default: 4)")
891     fetch_parser.add_argument("--retries", type=int, default=2, metavar="N",
892                                     help="Retries per video when a download fails or returns a
low format (default: 2)")
893     fetch_parser.add_argument("--retry-backoff", type=float, default=20.0,
metavar="SECONDS",
894                                     help="Base backoff between retries; grows linearly per
attempt (default: 20)")
895
896     # Convert-cvat command ? normalize a vendor CVAT export into our canonical CSV
897     cvat_parser = subparsers.add_parser(
898         "convert-cvat",
899         help="Convert a vendor CVAT rally export (bounding-box XML) into the canonical
per-rally CSV")
900     cvat_parser.add_argument("--input", required=True,
901                             help="Path to the CVAT XML export (image- or track-mode)")
902     cvat_parser.add_argument("--out", required=True,
903                             help="Path to write the canonical per-rally CSV (import-
local input)")
904     cvat_parser.add_argument("--video-path",
905                             help="Source video; its fps is read via ffprobe (seconds =
frame / fps)")
906     cvat_parser.add_argument("--fps", type=float,
907                             help="Override fps instead of probing the video")
908     cvat_parser.add_argument("--step", type=int,
909                             help="Frame sampling step (default: read from CVAT meta,
else inferred)")
910     cvat_parser.add_argument("--issues",
911                             help="Path to also write the human-readable issues report")
912
913     # Validate-delivery command ? run the full check suite + emit a structured report
914     vd_parser = subparsers.add_parser(
915         "validate-delivery",
916         help="Validate a rally delivery (CVAT export or canonical CSV) and emit report +
review sheet")
917     vd_parser.add_argument("--input", help="CVAT XML export (use with --video-
path/--fps)")
918     vd_parser.add_argument("--csv", help="An already-canonical per-rally CSV to
validate")
919     vd_parser.add_argument("--video-path",
920                             help="Source video; fps via ffprobe (when using --input)")
921     vd_parser.add_argument("--fps", type=float, help="Override fps (when using
--input)")
922     vd_parser.add_argument("--video-id", help="Names the report files (<id>_report.md,
etc.)")
923     vd_parser.add_argument("--out-dir", default="data/validation",
924                             help="Where to write the report + review sheet")
925
926     # Export command
927     export_parser = subparsers.add_parser(
928         "export", help="Export curated annotations as indexer-ready eval/training CSVs +
manifest")
929     export_parser.add_argument("--out-dir", default="data/eval_export",
930                             help="Output directory for per-video CSVs and
manifest.json")
931     export_parser.add_argument("--sport",

```

```

932             help="Restrict to a single sport (default: all
supported_sports)")
933     export_parser.add_argument("--include-staging", action="store_true",
934                               help="Also export videos still in staging (default:
curated only)")
935     export_parser.add_argument("--include-noncommercial", action="store_true",
936                               help="Include videos whose license is not commercially
safe (default: commercial only)")
937
938     args = parser.parse_args()
939     cli = SportsDataCLI()
940
941     if args.command == "status":
942         cli.status()
943     elif args.command == "curate":
944         cli.curate()
945     elif args.command == "import-local":
946         cli.import_local(args)
947     elif args.command == "convert-cvat":
948         cli.convert_cvat(args)
949     elif args.command == "validate-delivery":
950         cli.validate_delivery(args)
951     elif args.command == "export":
952         cli.export_eval_set(args)
953     elif args.command == "fetch":
954         cli.fetch(args)
955     else:
956         # Default: print help
957         parser.print_help()
958
959     if __name__ == "__main__":
960         main()
961

```

13. user (2026-06-10T06:00:24.906Z)

```

1     import sqlite3
2     import os
3     from contextlib import contextmanager
4     from typing import List, Dict, Any, Optional
5     from src.core.schemas.base import VideoMetadata, SportType, CurationStatus, LicenseType
6     from src.core.schemas.badminton import BadmintonRally
7     from src.core.schemas.tennis import TennisPoint
8     from src.core.schemas.cricket import CricketDelivery
9     from src.core.schemas.pickleball import PickleballRally
10    from src.core.schemas.tabletennis import TableTennisRally
11
12    # Maps each sport to its annotation table and the column(s) that order segments
13    # chronologically. All child tables expose (start_time, end_time), so a
14    # single generic reader can serve every sport ? adding a new sport is one entry
15    # here plus its CREATE TABLE, with no new query code.
16    SEGMENT_TABLES = {
17        SportType.BADMINTON: ("badminton_rallies", "rally_number"),
18        SportType.TENNIS: ("tennis_points", "point_number"),
19        SportType.CRICKET: ("cricket_deliveries", "over_num, ball_num"),
20        SportType.PICKLEBALL: ("pickleball_rallies", "rally_number"),
21        SportType.TABLE_TENNIS: ("tabletennis_rallies", "rally_number"),
22    }
23
24
25    class SportsDatabase:
26        def __init__(self, db_path: str):
27            self.db_path = db_path
28            # Ensure target directory exists
29            os.makedirs(os.path.dirname(os.path.abspath(db_path)), exist_ok=True)

```

```

30         self._init_db()
31
32     @contextmanager
33     def _get_connection(self):
34         """Yield a SQLite connection and always close it.
35
36         Used as ``with self._get_connection() as conn:``. A plain
37         ``sqlite3.connect(...)`` used as a context manager commits/rolls back the
38         transaction but does NOT close the connection, which leaks handles and
39         raises ResourceWarnings. This wrapper guarantees the connection is closed
40         on exit; writers still commit explicitly inside the block.
41         """
42         conn = sqlite3.connect(self.db_path)
43         conn.row_factory = sqlite3.Row
44         # SQLite disables foreign keys per-connection by default; enable so the
45         # ON DELETE CASCADE declarations on the child tables actually fire.
46         conn.execute("PRAGMA foreign_keys = ON")
47         try:
48             yield conn
49         finally:
50             conn.close()
51
52     def _init_db(self):
53         with self._get_connection() as conn:
54             cursor = conn.cursor()
55
56             # 1. Create common videos table
57             cursor.execute("""
58             CREATE TABLE IF NOT EXISTS videos (
59                 video_id TEXT PRIMARY KEY,
60                 sport TEXT NOT NULL,
61                 video_url TEXT,
62                 local_path TEXT,
63                 license_type TEXT NOT NULL,
64                 license_url TEXT,
65                 is_commercial INTEGER NOT NULL DEFAULT 0,
66                 status TEXT NOT NULL DEFAULT 'staging',
67                 match_name TEXT NOT NULL,
68                 fps REAL,
69                 resolution TEXT
70             )
71             """)
72
73             # 2. Create badminton_rallies table
74             cursor.execute("""
75             CREATE TABLE IF NOT EXISTS badminton_rallies (
76                 video_id TEXT NOT NULL,
77                 rally_number INTEGER NOT NULL,
78                 start_time REAL NOT NULL,
79                 end_time REAL NOT NULL,
80                 score_before TEXT,
81                 score_after TEXT,
82                 shots_count INTEGER,
83                 ending_reason TEXT,
84                 last_shot_outcome TEXT,
85                 PRIMARY KEY (video_id, rally_number),
86                 FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
87             )
88             """)
89
90             # 3. Create tennis_points table
91             cursor.execute("""
92             CREATE TABLE IF NOT EXISTS tennis_points (
93                 video_id TEXT NOT NULL,
94                 point_number INTEGER NOT NULL,

```

```

95         start_time REAL NOT NULL,
96         end_time REAL NOT NULL,
97         set_score TEXT,
98         game_score TEXT,
99         server TEXT,
100        ending_type TEXT,
101        PRIMARY KEY (video_id, point_number),
102        FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
103    )
104    """)
105
106    # 4. Create cricket_deliveries table
107    cursor.execute("""
108    CREATE TABLE IF NOT EXISTS cricket_deliveries (
109        video_id TEXT NOT NULL,
110        over_num INTEGER NOT NULL,
111        ball_num INTEGER NOT NULL,
112        start_time REAL NOT NULL,
113        end_time REAL NOT NULL,
114        batsman TEXT,
115        bowler TEXT,
116        runs INTEGER,
117        wicket INTEGER DEFAULT 0,
118        extras INTEGER DEFAULT 0,
119        PRIMARY KEY (video_id, over_num, ball_num),
120        FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
121    )
122    """)
123
124    # 5. Create pickleball_rallies table
125    cursor.execute("""
126    CREATE TABLE IF NOT EXISTS pickleball_rallies (
127        video_id TEXT NOT NULL,
128        rally_number INTEGER NOT NULL,
129        start_time REAL NOT NULL,
130        end_time REAL NOT NULL,
131        score_before TEXT,
132        score_after TEXT,
133        shots_count INTEGER,
134        ending_reason TEXT,
135        last_shot_outcome TEXT,
136        PRIMARY KEY (video_id, rally_number),
137        FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
138    )
139    """)
140
141    # 6. Create tabletennis_rallies table
142    cursor.execute("""
143    CREATE TABLE IF NOT EXISTS tabletennis_rallies (
144        video_id TEXT NOT NULL,
145        rally_number INTEGER NOT NULL,
146        start_time REAL NOT NULL,
147        end_time REAL NOT NULL,
148        score_before TEXT,
149        score_after TEXT,
150        shots_count INTEGER,
151        ending_reason TEXT,
152        last_shot_outcome TEXT,
153        PRIMARY KEY (video_id, rally_number),
154        FOREIGN KEY (video_id) REFERENCES videos(video_id) ON DELETE CASCADE
155    )
156    """)
157
158    # Migration: add last_shot_outcome to rally tables that predate the
159    # column. CREATE TABLE IF NOT EXISTS won't alter an existing table, so

```



```

160         # do it explicitly and idempotently (skip if the column is already there).
161         for tbl in ("badminton_rallies", "pickleball_rallies",
"tabletennis_rallies"):
162             existing = [r[1] for r in cursor.execute(f"PRAGMA
table_info({tbl})).fetchall()]
163             if "last_shot_outcome" not in existing:
164                 cursor.execute(f"ALTER TABLE {tbl} ADD COLUMN last_shot_outcome
TEXT")
165
166             conn.commit()
167
168     def insert_video(self, video: VideoMetadata) -> bool:
169         """Insert a video entry, or update it in place if it already exists.
170
171         Uses an UPSERT rather than INSERT OR REPLACE: REPLACE would DELETE the
172         existing row, which (with foreign keys enabled) would cascade-delete the
173         video's child annotations. UPSERT mutates the row in place and leaves
174         children untouched.
175         """
176         with self._get_connection() as conn:
177             cursor = conn.cursor()
178             cursor.execute("""
179             INSERT INTO videos (
180                 video_id, sport, video_url, local_path, license_type,
181                 license_url, is_commercial, status, match_name, fps, resolution
182             ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
183             ON CONFLICT(video_id) DO UPDATE SET
184                 sport=excluded.sport,
185                 video_url=excluded.video_url,
186                 local_path=excluded.local_path,
187                 license_type=excluded.license_type,
188                 license_url=excluded.license_url,
189                 is_commercial=excluded.is_commercial,
190                 status=excluded.status,
191                 match_name=excluded.match_name,
192                 fps=excluded.fps,
193                 resolution=excluded.resolution
194             """, (
195                 video.video_id,
196                 video.sport.value,
197                 video.video_url,
198                 video.local_path,
199                 video.license_type.value,
200                 video.license_url,
201                 1 if video.is_commercial else 0,
202                 video.status.value,
203                 video.match_name,
204                 video.fps,
205                 video.resolution
206             ))
207             conn.commit()
208             return True
209
210     def insert_badminton_rallies(self, rallies: List[BadmintonRally]) -> int:
211         if not rallies:
212             return 0
213         with self._get_connection() as conn:
214             cursor = conn.cursor()
215             # First clear existing rallies for this video_id to avoid key conflicts on
updates
216             video_id = rallies[0].video_id
217             cursor.execute("DELETE FROM badminton_rallies WHERE video_id = ?",
(video_id,))
218
219             for rally in rallies:

```

```

220         cursor.execute("""
221         INSERT INTO badminton_rallies (
222             video_id, rally_number, start_time, end_time,
223             score_before, score_after, shots_count, ending_reason,
last_shot_outcome
224         ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)
225         """, (
226             rally.video_id,
227             rally.rally_number,
228             rally.start_time,
229             rally.end_time,
230             rally.score_before,
231             rally.score_after,
232             rally.shots_count,
233             rally.ending_reason,
234             rally.last_shot_outcome
235         ))
236         conn.commit()
237         return len(rallies)
238
239     def insert_tennis_points(self, points: List[TennisPoint]) -> int:
240         if not points:
241             return 0
242         with self._get_connection() as conn:
243             cursor = conn.cursor()
244             video_id = points[0].video_id
245             cursor.execute("DELETE FROM tennis_points WHERE video_id = ?", (video_id,))
246
247             for pt in points:
248                 cursor.execute("""
249                 INSERT INTO tennis_points (
250                     video_id, point_number, start_time, end_time,
251                     set_score, game_score, server, ending_type
252                 ) VALUES (?, ?, ?, ?, ?, ?, ?, ?)
253                 """, (
254                     pt.video_id,
255                     pt.point_number,
256                     pt.start_time,
257                     pt.end_time,
258                     pt.set_score,
259                     pt.game_score,
260                     pt.server,
261                     pt.ending_type
262                 ))
263                 conn.commit()
264             return len(points)
265
266     def insert_cricket_deliveries(self, deliveries: List[CricketDelivery]) -> int:
267         if not deliveries:
268             return 0
269         with self._get_connection() as conn:
270             cursor = conn.cursor()
271             video_id = deliveries[0].video_id
272             cursor.execute("DELETE FROM cricket_deliveries WHERE video_id = ?",
(video_id,))
273
274             for deliv in deliveries:
275                 cursor.execute("""
276                 INSERT INTO cricket_deliveries (
277                     video_id, over_num, ball_num, start_time, end_time,
278                     batsman, bowler, runs, wicket, extras
279                 ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
280                 """, (
281                     deliv.video_id,
282                     deliv.over,

```

```

283         deliv.ball,
284         deliv.start_time,
285         deliv.end_time,
286         deliv.batsman,
287         deliv.bowler,
288         deliv.runs,
289         1 if deliv.wicket else 0,
290         deliv.extras
291     ))
292     conn.commit()
293     return len(deliveries)
294
295     def insert_pickleball_rallies(self, rallies: List[PickleballRally]) -> int:
296         if not rallies:
297             return 0
298         with self._get_connection() as conn:
299             cursor = conn.cursor()
300             video_id = rallies[0].video_id
301             cursor.execute("DELETE FROM pickleball_rallies WHERE video_id = ?",
302 (video_id,))
303             for rally in rallies:
304                 cursor.execute("""
305                     INSERT INTO pickleball_rallies (
306                         video_id, rally_number, start_time, end_time,
307                         score_before, score_after, shots_count, ending_reason,
308                         last_shot_outcome
309                     ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)
310                     """, (
311                         rally.video_id,
312                         rally.rally_number,
313                         rally.start_time,
314                         rally.end_time,
315                         rally.score_before,
316                         rally.score_after,
317                         rally.shots_count,
318                         rally.ending_reason,
319                         rally.last_shot_outcome
320                     ))
321             conn.commit()
322             return len(rallies)
323
324     def insert_tabletennis_rallies(self, rallies: List[TableTennisRally]) -> int:
325         if not rallies:
326             return 0
327         with self._get_connection() as conn:
328             cursor = conn.cursor()
329             video_id = rallies[0].video_id
330             cursor.execute("DELETE FROM tabletennis_rallies WHERE video_id = ?",
331 (video_id,))
332             for rally in rallies:
333                 cursor.execute("""
334                     INSERT INTO tabletennis_rallies (
335                         video_id, rally_number, start_time, end_time,
336                         score_before, score_after, shots_count, ending_reason,
337                         last_shot_outcome
338                     ) VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)
339                     """, (
340                         rally.video_id,
341                         rally.rally_number,
342                         rally.start_time,
343                         rally.end_time,
344                         rally.score_before,
345                         rally.score_after,

```

```

344         rally.shots_count,
345         rally.ending_reason,
346         rally.last_shot_outcome
347     ))
348     conn.commit()
349     return len(rallies)
350
351 def get_video(self, video_id: str) -> Optional[VideoMetadata]:
352     with self._get_connection() as conn:
353         cursor = conn.cursor()
354         cursor.execute("SELECT * FROM videos WHERE video_id = ?", (video_id,))
355         row = cursor.fetchone()
356         if row:
357             return VideoMetadata(
358                 video_id=row['video_id'],
359                 sport=SportType(row['sport']),
360                 video_url=row['video_url'],
361                 local_path=row['local_path'],
362                 license_type=LicenseType(row['license_type']),
363                 license_url=row['license_url'],
364                 is_commercial=bool(row['is_commercial']),
365                 status=CurationStatus(row['status']),
366                 match_name=row['match_name'],
367                 fps=row['fps'],
368                 resolution=row['resolution']
369             )
370         return None
371
372 def get_videos_by_status(self, status: CurationStatus, sport: Optional[SportType] =
None) -> List[VideoMetadata]:
373     query = "SELECT * FROM videos WHERE status = ?"
374     params = [status.value]
375     if sport:
376         query += " AND sport = ?"
377         params.append(sport.value)
378
379     videos = []
380     with self._get_connection() as conn:
381         cursor = conn.cursor()
382         cursor.execute(query, params)
383         for row in cursor.fetchall():
384             videos.append(VideoMetadata(
385                 video_id=row['video_id'],
386                 sport=SportType(row['sport']),
387                 video_url=row['video_url'],
388                 local_path=row['local_path'],
389                 license_type=LicenseType(row['license_type']),
390                 license_url=row['license_url'],
391                 is_commercial=bool(row['is_commercial']),
392                 status=CurationStatus(row['status']),
393                 match_name=row['match_name'],
394                 fps=row['fps'],
395                 resolution=row['resolution']
396             ))
397     return videos
398
399 def get_status_counts(self) -> List[tuple]:
400     """Return (sport, status, count) rows for the status summary.
401
402     Public read API so callers (e.g. the CLI `status` command) don't have to
403     reach into the private `_get_connection` context manager.
404     """
405     with self._get_connection() as conn:
406         cursor = conn.cursor()
407         cursor.execute(

```

```

408         "SELECT sport, status, COUNT(*) as cnt FROM videos GROUP BY status,
sport"
409     )
410     return [(r["sport"], r["status"], r["cnt"]) for r in cursor.fetchall()]
411
412 def get_badminton_rallies(self, video_id: str) -> List[BadmintonRally]:
413     rallies = []
414     with self._get_connection() as conn:
415         cursor = conn.cursor()
416         cursor.execute("""
417         SELECT * FROM badminton_rallies
418         WHERE video_id = ?
419         ORDER BY rally_number
420         """, (video_id,))
421         for row in cursor.fetchall():
422             rallies.append(BadmintonRally(
423                 video_id=row['video_id'],
424                 rally_number=row['rally_number'],
425                 start_time=row['start_time'],
426                 end_time=row['end_time'],
427                 score_before=row['score_before'],
428                 score_after=row['score_after'],
429                 shots_count=row['shots_count'],
430                 ending_reason=row['ending_reason'],
431                 last_shot_outcome=row['last_shot_outcome']
432             ))
433     return rallies
434
435 def get_tennis_points(self, video_id: str) -> List[TennisPoint]:
436     points = []
437     with self._get_connection() as conn:
438         cursor = conn.cursor()
439         cursor.execute("""
440         SELECT * FROM tennis_points
441         WHERE video_id = ?
442         ORDER BY point_number
443         """, (video_id,))
444         for row in cursor.fetchall():
445             points.append(TennisPoint(
446                 video_id=row['video_id'],
447                 point_number=row['point_number'],
448                 start_time=row['start_time'],
449                 end_time=row['end_time'],
450                 set_score=row['set_score'],
451                 game_score=row['game_score'],
452                 server=row['server'],
453                 ending_type=row['ending_type']
454             ))
455     return points
456
457 def get_cricket_deliveries(self, video_id: str) -> List[CricketDelivery]:
458     deliveries = []
459     with self._get_connection() as conn:
460         cursor = conn.cursor()
461         cursor.execute("""
462         SELECT * FROM cricket_deliveries
463         WHERE video_id = ?
464         ORDER BY over_num, ball_num
465         """, (video_id,))
466         for row in cursor.fetchall():
467             deliveries.append(CricketDelivery(
468                 video_id=row['video_id'],
469                 over=row['over_num'],
470                 ball=row['ball_num'],
471                 start_time=row['start_time'],

```

```

472         end_time=row['end_time'],
473         batsman=row['batsman'],
474         bowler=row['bowler'],
475         runs=row['runs'],
476         wicket=bool(row['wicket']),
477         extras=row['extras']
478     ))
479     return deliveries
480
481 def get_pickleball_rallies(self, video_id: str) -> List[PickleballRally]:
482     rallies = []
483     with self._get_connection() as conn:
484         cursor = conn.cursor()
485         cursor.execute("""
486         SELECT * FROM pickleball_rallies
487         WHERE video_id = ?
488         ORDER BY rally_number
489         """, (video_id,))
490         for row in cursor.fetchall():
491             rallies.append(PickleballRally(
492                 video_id=row['video_id'],
493                 rally_number=row['rally_number'],
494                 start_time=row['start_time'],
495                 end_time=row['end_time'],
496                 score_before=row['score_before'],
497                 score_after=row['score_after'],
498                 shots_count=row['shots_count'],
499                 ending_reason=row['ending_reason'],
500                 last_shot_outcome=row['last_shot_outcome']
501             ))
502     return rallies
503
504 def get_tabletennis_rallies(self, video_id: str) -> List[TableTennisRally]:
505     rallies = []
506     with self._get_connection() as conn:
507         cursor = conn.cursor()
508         cursor.execute("""
509         SELECT * FROM tabletennis_rallies
510         WHERE video_id = ?
511         ORDER BY rally_number
512         """, (video_id,))
513         for row in cursor.fetchall():
514             rallies.append(TableTennisRally(
515                 video_id=row['video_id'],
516                 rally_number=row['rally_number'],
517                 start_time=row['start_time'],
518                 end_time=row['end_time'],
519                 score_before=row['score_before'],
520                 score_after=row['score_after'],
521                 shots_count=row['shots_count'],
522                 ending_reason=row['ending_reason'],
523                 last_shot_outcome=row['last_shot_outcome']
524             ))
525     return rallies
526
527 def get_segment_intervals(self, video_id: str, sport: SportType) -> List[tuple]:
528     """Return chronologically-ordered (start_time, end_time) tuples for a video,
529     regardless of sport. This is the sport-agnostic view used to export
530     indexer-ready evaluation/training labels.
531
532     Sports that have no annotation table yet (e.g. tabletennis, pickleball)
533     yield an empty list rather than raising.
534     """
535     mapping = SEGMENT_TABLES.get(sport)
536     if mapping is None:

```

```

537         return []
538         table, order_by = mapping # both come from the trusted map above, never user
input
539         with self._get_connection() as conn:
540             cursor = conn.cursor()
541             cursor.execute(
542                 f"SELECT start_time, end_time FROM {table} "
543                 f"WHERE video_id = ? ORDER BY {order_by}",
544                 (video_id,),
545             )
546             return [(row["start_time"], row["end_time"]) for row in cursor.fetchall()]
547
548     def update_status(self, video_id: str, status: CurationStatus) -> bool:
549         with self._get_connection() as conn:
550             cursor = conn.cursor()
551             cursor.execute("UPDATE videos SET status = ? WHERE video_id = ?",
(status.value, video_id))
552             conn.commit()
553             return cursor.rowcount > 0
554
555     def delete_video(self, video_id: str) -> bool:
556         """Deletes a video and cascading child entries."""
557         with self._get_connection() as conn:
558             cursor = conn.cursor()
559             # Execute cascade deletes manually just in case SQLite foreign key support
is not active
560             cursor.execute("DELETE FROM badminton_rallies WHERE video_id = ?",
(video_id,))
561             cursor.execute("DELETE FROM tennis_points WHERE video_id = ?", (video_id,))
562             cursor.execute("DELETE FROM cricket_deliveries WHERE video_id = ?",
(video_id,))
563             cursor.execute("DELETE FROM pickleball_rallies WHERE video_id = ?",
(video_id,))
564             cursor.execute("DELETE FROM tabletennis_rallies WHERE video_id = ?",
(video_id,))
565             cursor.execute("DELETE FROM videos WHERE video_id = ?", (video_id,))
566             conn.commit()
567             return cursor.rowcount > 0
568

```

14. user (2026-06-10T06:00:33.148Z)

```

1     """Adapter: CVAT rally exports -> our canonical per-rally records.
2
3     Why this exists
4     -----
5     Our annotation brief (`docs/annotation/ROORA_BRIEF.md`) asks vendors for a
6     per-rally CSV that imports directly via curate import-local. A vendor may
7     instead do the work in CVAT and hand back the CVAT export: per-frame bounding
8     boxes, with the rally fields (rally_number, shots_count,
9     ending_reason ...) carried as CVAT object attributes on each box. The box
10    geometry is irrelevant to us ? it is just a carrier the rater attached the rally
11    metadata to.
12
13    This module is the ingestion-side adapter that absorbs that vendor format so the
14    modeling repo never has to. It does two things, both deliberately on our side
15    so we fix mechanical problems rather than renegotiating them:
16
17    1. Recompute time from the frame range. A rally's interval is taken from the
18    first/last frame a box for that rally appears on:
19    seconds = frame / fps at the video's true fps. We ignore any
20    start_time/end_time attributes the vendor wrote ? a CVAT job that
21    samples every step frames and assumes 30 fps yields systematically wrong
22    times (we have seen start_time == frame / (step * 30), i.e. off by the
23    sampling step at the wrong fps). The frame index is reliable; the vendor's

```

```

24         seconds are not.
25
26     2. **Surface ? never silently fix ? content defects.** Overlapping rallies,
27         zero-duration rallies, ending_reason values outside the agreed set, and
28         too-coarse sampling are reported as :class:`Issue`s. Guessing a corrected
29         rally boundary would poison ground truth, so that stays a human/vendor call.
30
31     Supported input: the single-file CVAT XML export, in either **image mode**
32     (``<image><box>...``) or **track/interpolation mode** (``<track><box frame=...>``).
33     The two PASCAL-VOC / COCO sibling exports encode the same data and can be added
34     later if a vendor sends only those.
35     """
36     from __future__ import annotations
37
38     import os
39     import xml.etree.ElementTree as ET
40     from collections import Counter
41     from dataclasses import dataclass, field
42     from typing import Any, Dict, List, Optional, Tuple
43
44     from src.core.parsing import safe_int
45     from src.core.downloader import probe_resolution_fps
46
47     # The agreed cross-sport ending_reason vocabulary for net-separated sports.
48     # Kept here as a soft check only (a warning, never a hard failure) so this
49     # adapter does not depend on the annotation-spec wording.
50     DEFAULT_ALLOWED_REASONS = {
51         "winner",
52         "forced_error",
53         "unforced_error",
54         "service_fault",
55         "let",
56         "other",
57     }
58
59     # The observational outcome of the last shot ? orthogonal to ending_reason
60     # (winner can be 'in'; an error can be 'net' or 'out'; service_fault/let -> 'n_a').
61     DEFAULT_ALLOWED_OUTCOMES = {"in", "net", "out", "n_a"}
62
63     # CVAT attribute names we map onto canonical columns. The vendor's CVAT label
64     # schema has used suffixed names like "score_before (A-B)"; we match leniently.
65     _ATTR_ALIASES = {
66         "rally_number": ("rally_number",),
67         "shots_count": ("shots_count",),
68         "ending_reason": ("ending_reason",),
69         "last_shot_outcome": ("last_shot_outcome", "outcome", "last_shot_outcome
70 (in/net/out)"),
71         "score_before": ("score_before", "score_before (a-b)"),
72         "score_after": ("score_after", "score_after (a-b)"),
73         "notes": ("notes", "note"),
74     }
75
76     CANONICAL_COLUMNS = [
77         "rally_number", "start_mmss", "end_mmss", "start_time", "end_time",
78         "shots_count", "ending_reason", "last_shot_outcome", "score_before", "score_after",
79         "notes",
80     ]
81
82     @dataclass
83     class RallyRecord:
84         """One rally recovered from the CVAT export, time re-derived from frames."""
85         rally_number: int
86         start_time: float          # seconds, = frame_min / fps
87         end_time: float            # seconds, = frame_max / fps (see padding note)

```



```

87     shots_count: Optional[int]
88     ending_reason: Optional[str]
89     last_shot_outcome: Optional[str]
90     score_before: Optional[str]
91     score_after: Optional[str]
92     notes: Optional[str]
93     frame_min: int
94     frame_max: int
95     n_boxes: int
96     vendor_start_time: Optional[str] = None    # the (discarded) attribute value
97     vendor_end_time: Optional[str] = None
98
99     @property
100     def duration(self) -> float:
101         return self.end_time - self.start_time
102
103
104 @dataclass
105 class Issue:
106     """A defect found in the delivery. ``blocker`` issues stop a clean import."""
107     severity: str    # 'blocker' | 'high' | 'medium' | 'low'
108     code: str
109     message: str
110     rally_number: Optional[int] = None
111
112
113 @dataclass
114 class ConversionResult:
115     records: List[RallyRecord]
116     issues: List[Issue]
117     fps: float
118     step: int
119     meta: Dict[str, Any] = field(default_factory=dict)
120
121     @property
122     def n_blockers(self) -> int:
123         return sum(1 for i in self.issues if i.severity == "blocker")
124
125
126 def _mmss(seconds: float) -> str:
127     """Format seconds as ``m:ss.mmm`` (matches the brief's mm:ss columns)."""
128     m = int(seconds // 60)
129     s = seconds - m * 60
130     return f"{m}:{s:06.3f}"
131
132
133 def _attr_lookup(attrs: Dict[str, str], canonical: str) -> Optional[str]:
134     for alias in _ATTR_ALIASES[canonical]:
135         if alias in attrs and attrs[alias] != "":
136             return attrs[alias]
137     return None
138
139
140 def parse_cvat_boxes(xml_path: str) -> Tuple[Dict[str, Any], List[Dict[str, Any]]]:
141     """Parse a CVAT XML export into (meta, boxes).
142
143     ``boxes`` is a flat list of ``{"frame": int, "attrs": {name: value}}`` ? one
144     entry per visible box, from either image-mode or track-mode exports.
145     ``meta`` carries ``step`` (frame_filter) and ``stop_frame`` when present.
146     """
147     if not os.path.exists(xml_path):
148         raise FileNotFoundError(f"CVAT export not found: {xml_path}")
149
150     root = ET.parse(xml_path).getroot()
151     meta: Dict[str, Any] = {}

```

```

152
153 job = root.find("./job")
154 if job is not None:
155     stop = job.findtext("stop_frame")
156     if stop is not None and stop.strip().isdigit():
157         meta["stop_frame"] = int(stop.strip())
158     ff = job.findtext("frame_filter") or ""
159     if "step=" in ff:
160         try:
161             meta["step"] = int(ff.split("step=")[1].split()[0])
162         except (ValueError, IndexError):
163             pass
164
165 def _attrs(el: ET.Element) -> Dict[str, str]:
166     out: Dict[str, str] = {}
167     for a in el.findall("attribute"):
168         out[(a.get("name") or "").strip().lower()] = (a.text or "").strip()
169     return out
170
171 boxes: List[Dict[str, Any]] = []
172
173 # Image mode: <image id=frame> ... <box>...</box>
174 for img in root.findall("image"):
175     try:
176         frame = int(img.get("id"))
177     except (TypeError, ValueError):
178         continue
179     for box in img.findall("box"):
180         boxes.append({"frame": frame, "attrs": _attrs(box)})
181
182 # Track mode: <track> ... <box frame=N outside=0>...</box>
183 for track in root.findall("track"):
184     for box in track.findall("box"):
185         if box.get("outside") == "1":
186             continue # a track's terminating keyframe carries no real label
187         try:
188             frame = int(box.get("frame"))
189         except (TypeError, ValueError):
190             continue
191         boxes.append({"frame": frame, "attrs": _attrs(box)})
192
193 return meta, boxes
194
195
196 def _infer_step(frames: List[int]) -> int:
197     """Most common positive gap between consecutive distinct sampled frames."""
198     uniq = sorted(set(frames))
199     gaps = [b - a for a, b in zip(uniq, uniq[1:]) if b - a > 0]
200     if not gaps:
201         return 1
202     return Counter(gaps).most_common(1)[0][0]
203
204
205 def build_records(boxes: List[Dict[str, Any]], fps: float) -> List[RallyRecord]:
206     """Group boxes by rally_number and derive each rally's frame range -> seconds."""
207     if fps is None or fps <= 0:
208         raise ValueError("fps must be a positive number to convert frames to seconds")
209
210     by_rally: Dict[int, List[Dict[str, Any]]] = {}
211     for b in boxes:
212         rn = safe_int(_attr_lookup(b["attrs"], "rally_number"), None)
213         if rn is None:
214             continue
215         by_rally.setdefault(rn, []).append(b)
216

```

```

217 records: List[RallyRecord] = []
218 for rn in sorted(by_rally):
219     items = by_rally[rn]
220     frames = sorted(b["frame"] for b in items)
221     attrs = items[0]["attrs"]
222     fmin, fmax = frames[0], frames[-1]
223     records.append(RallyRecord(
224         rally_number=rn,
225         start_time=fmin / fps,
226         end_time=fmax / fps,
227         shots_count=safe_int(_attr_lookup(attrs, "shots_count"), None),
228         ending_reason=_attr_lookup(attrs, "ending_reason"),
229         last_shot_outcome=_attr_lookup(attrs, "last_shot_outcome"),
230         score_before=_attr_lookup(attrs, "score_before"),
231         score_after=_attr_lookup(attrs, "score_after"),
232         notes=_attr_lookup(attrs, "notes"),
233         frame_min=fmin,
234         frame_max=fmax,
235         n_boxes=len(frames),
236         vendor_start_time=attrs.get("start_time"),
237         vendor_end_time=attrs.get("end_time"),
238     ))
239 return records
240
241
242 def detect_issues(records: List[RallyRecord],
243                 stop_frame: Optional[int] = None,
244                 allowed_reasons: Optional[set] = None,
245                 allowed_outcomes: Optional[set] = None,
246                 max_gap_s: float = 15.0) -> List[Issue]:
247     """Find the defects that block or degrade a clean ground-truth import.
248
249     Note: we deliberately do NOT flag the sampling density (step/fps). A boundary
250     slack of ~0.5s is acceptable for this pipeline because the downstream
251     highlight compiler pads every clip on output (badminton-highlight-indexer,
252     backend/stitcher/compiler.py: begin_padding=0.5s, end_padding=1.0s), which
253     absorbs the sampling uncertainty when highlights are trimmed.
254     """
255     allowed = allowed_reasons if allowed_reasons is not None else
256     DEFAULT_ALLOWED_REASONS
257     allowed_out = allowed_outcomes if allowed_outcomes is not None else
258     DEFAULT_ALLOWED_OUTCOMES
259     issues: List[Issue] = []
260
261     ordered = sorted(records, key=lambda r: r.rally_number)
262     for i, r in enumerate(ordered):
263         if r.end_time <= r.start_time:
264             issues.append(Issue(
265                 "blocker", "zero_duration",
266                 f"Rally {r.rally_number}: end <= start "
267                 f"({frames {r.frame_min}-{r.frame_max}}); single-sample or stray box.",
268                 r.rally_number))
269
270         if r.ending_reason is not None and r.ending_reason not in allowed:
271             issues.append(Issue(
272                 "medium", "ending_reason_vocab",
273                 f"Rally {r.rally_number}: ending_reason '{r.ending_reason}' "
274                 f"not in the agreed set {sorted(allowed)}.",
275                 r.rally_number))
276
277         if r.last_shot_outcome is not None and r.last_shot_outcome not in allowed_out:
278             issues.append(Issue(
279                 "low", "outcome_vocab",
280                 f"Rally {r.rally_number}: last_shot_outcome '{r.last_shot_outcome}' "
281                 f"not in {sorted(allowed_out)}.",

```

```

280         r.rally_number))
281
282     # Inflated span: a long rally with implausibly few shots usually means a
283     # stray/duplicated box stretched the frame range past the real ending.
284     if r.shots_count and r.duration >= 5.0:
285         sps = r.shots_count / r.duration
286         if sps < 0.3:
287             issues.append(Issue(
288                 "medium", "shot_density",
289                 f"Rally {r.rally_number}: {r.shots_count} shots over "
290                 f"{r.duration:.1f}s ({sps:.2f}/s) is implausibly low ? the span "
291                 f"may be inflated by a stray box.",
292                 r.rally_number))
293
294     if stop_frame is not None and r.frame_max > stop_frame:
295         issues.append(Issue(
296             "high", "frame_out_of_range",
297             f"Rally {r.rally_number}: frame {r.frame_max} exceeds the video's "
298             f"last frame {stop_frame}.",
299             r.rally_number))
300
301     if i > 0:
302         prev = ordered[i - 1]
303         if r.start_time < prev.end_time:
304             issues.append(Issue(
305                 "blocker", "overlap",
306                 f"Rally {prev.rally_number} (ends {prev.end_time:.2f}s) overlaps "
307                 f"Rally {r.rally_number} (starts {r.start_time:.2f}s) by "
308                 f"{prev.end_time - r.start_time:.2f}s ? rallies must be disjoint.",
309                 r.rally_number))
310             elif r.start_time - prev.end_time > max_gap_s:
311                 # A long unlabeled stretch between two rallies is a cheap, label-only
312                 # second opinion for the worst error (a missed rally). It won't catch a
313                 # rally MERGED into its neighbour's span (shot_density/inflated catches
314                 # that), but it does catch an omitted rally that left a visible hole.
315                 gap = r.start_time - prev.end_time
316                 issues.append(Issue(
317                     "medium", "possible_missed_rally",
318                     f"{gap:.1f}s of unlabeled time between Rally {prev.rally_number} "
319                     f"(ends {prev.end_time:.1f}s) and Rally {r.rally_number} "
320                     f"(starts {r.start_time:.1f}s) ? check for an unlabeled rally in the
321 gap.",
322                     r.rally_number))
323
324     return issues
325
326 def convert(xml_path: str,
327            fps: Optional[float] = None,
328            video_path: Optional[str] = None,
329            step: Optional[int] = None,
330            allowed_reasons: Optional[set] = None) -> ConversionResult:
331     """Full pipeline: parse CVAT XML -> records (+ issues) at the true fps.
332
333     ``fps`` resolution order: explicit ``fps`` arg > ffprobe of ``video_path`` >
334     the CVAT meta (none) -> error. We never fall back to a guessed 30 fps, since
335     that silent guess is exactly the bug this adapter exists to undo.
336     """
337     meta, boxes = parse_cvat_boxes(xml_path)
338
339     if fps is None and video_path:
340         _, _, probed = probe_resolution_fps(video_path)
341         fps = probed
342
343     if fps is None or fps <= 0:
344         raise ValueError(
345             "Could not determine fps. Pass fps=... or a readable video_path so "

```

```

344         "frames can be converted to seconds (seconds = frame / fps).")
345
346     if step is None:
347         step = meta.get("step") or _infer_step([b["frame"] for b in boxes])
348
349     records = build_records(boxes, fps)
350     issues = detect_issues(records,
351                             stop_frame=meta.get("stop_frame"),
352                             allowed_reasons=allowed_reasons)
353     return ConversionResult(records=records, issues=issues, fps=fps, step=step,
354                             meta=meta)
355
356 def write_canonical_csv(records: List[RallyRecord], out_path: str) -> None:
357     """Write records as the canonical per-rally CSV that ``import-local`` reads."""
358     import csv
359     os.makedirs(os.path.dirname(os.path.abspath(out_path)), exist_ok=True)
360     with open(out_path, "w", newline="", encoding="utf-8") as f:
361         w = csv.writer(f)
362         w.writerow(CANONICAL_COLUMNS)
363         for r in sorted(records, key=lambda x: x.rally_number):
364             w.writerow([
365                 r.rally_number,
366                 _mmss(r.start_time), _mmss(r.end_time),
367                 f"{r.start_time:.3f}", f"{r.end_time:.3f}",
368                 r.shots_count if r.shots_count is not None else "",
369                 r.ending_reason or "",
370                 r.last_shot_outcome or "",
371                 r.score_before or "", r.score_after or "", r.notes or "",
372             ])
373
374
375 def render_issue_report(result: ConversionResult, source: str) -> str:
376     """A human-readable issue report (for review / sending back to the vendor)."""
377     lines = [
378         f"CVAT rally conversion report ? {source}",
379         f"  fps used          : {result.fps:.3f}",
380         f"  sampling step      : {result.step} frames (~{result.step / result.fps:.2f}s)",
381         f"  rallies parsed     : {len(result.records)}",
382         f"  issues              : {len(result.issues)} ({result.n_blockers} blocker)",
383         ""
384     ]
385     if not result.issues:
386         lines.append("No issues detected.")
387     else:
388         for sev in ("blocker", "high", "medium", "low"):
389             group = [i for i in result.issues if i.severity == sev]
390             if not group:
391                 continue
392             lines.append(f"[{sev.upper()}]")
393             for i in group:
394                 lines.append(f"  - ({i.code}) {i.message}")
395             lines.append("")
396     return "\n".join(lines)
397
398
399 # ----- #
400 # Validation reporting ? used by `curate validate-delivery`. #
401 # These turn a (records, issues) pair into the artifacts a human reviews: #
402 # a machine summary (JSON), a human report (md), and an annotatable sheet. #
403 # ----- #
404
405 def _parse_secs(value: Optional[str]) -> Optional[float]:
406     """Decimal seconds or mm:ss / hh:mm:ss -> float; None if blank/unparseable."""
407     s = (value or "").strip()

```

```

408         if not s:
409             return None
410         if ":" in s:
411             try:
412                 total = 0.0
413                 for part in s.split(":"):
414                     total = total * 60 + float(part)
415                 return total
416             except ValueError:
417                 return None
418         try:
419             return float(s)
420         except ValueError:
421             return None
422
423
424 def load_canonical_csv(csv_path: str) -> List[RallyRecord]:
425     """Load a canonical per-rally CSV into RallyRecords, so a *CSV* delivery (not
426     just a CVAT export) can be validated. Decimal start/end preferred, mm:ss
427     fallback; rows with no usable start/end are skipped. (frame_* are -1 here, as a
428     CSV carries no frame info ? only the frame-derived checks become inapplicable.)
429     """
430     import csv as _csv
431     if not os.path.exists(csv_path):
432         raise FileNotFoundError(f"CSV not found: {csv_path}")
433     records: List[RallyRecord] = []
434     with open(csv_path, newline="", encoding="utf-8") as f:
435         reader = _csv.DictReader(f)
436         reader.fieldnames = [(n or "").strip().lower() for n in (reader.fieldnames or
437 [])]
438         for raw in reader:
439             row = {(k or "").strip().lower(): v for k, v in raw.items()}
440             rn = safe_int(row.get("rally_number"), None)
441             if rn is None:
442                 continue
443             start = _parse_secs(row.get("start_time"))
444             if start is None:
445                 start = _parse_secs(row.get("start_mmss"))
446             end = _parse_secs(row.get("end_time"))
447             if end is None:
448                 end = _parse_secs(row.get("end_mmss"))
449             if start is None or end is None:
450                 continue
451             records.append(RallyRecord(
452                 rally_number=rn, start_time=start, end_time=end,
453                 shots_count=safe_int(row.get("shots_count"), None),
454                 ending_reason=(row.get("ending_reason") or "").strip() or None,
455                 last_shot_outcome=(row.get("last_shot_outcome") or "").strip() or None,
456                 score_before=(row.get("score_before") or "").strip() or None,
457                 score_after=(row.get("score_after") or "").strip() or None,
458                 notes=(row.get("notes") or "").strip() or None,
459                 frame_min=-1, frame_max=-1, n_boxes=0,
460             ))
461         return records
462
463 def build_summary(records: List[RallyRecord], issues: List[Issue],
464                 source: str = "", fps: Optional[float] = None) -> Dict[str, Any]:
465     """Machine-readable validation summary (JSON / CI)."""
466     by_sev = {s: [i for i in issues if i.severity == s]
467               for s in ("blocker", "high", "medium", "low")}
468     flagged = sorted([i.rally_number for i in issues if i.rally_number is not None])
469     return {
470         "source": source,
471         "fps": fps,

```

```

472         "n_rallies": len(records),
473         "n_issues": len(issues),
474         "counts_by_severity": {s: len(v) for s, v in by_sev.items()},
475         "n_blockers": len(by_sev["blocker"]),
476         "passed": len(by_sev["blocker"]) == 0,
477         "flagged_rallies": flagged,
478         "issues": [{"severity": i.severity, "code": i.code,
479                     "rally_number": i.rally_number, "message": i.message} for i in
issues],
480     }
481
482
483     def render_validation_md(records: List[RallyRecord], issues: List[Issue], source: str)
-> str:
484         """Human validation report: scoreboard + issues by severity + rallies to review."""
485         n = len(records)
486         by_sev = {s: [i for i in issues if i.severity == s]
487                   for s in ("blocker", "high", "medium", "low")}
488         flagged = sorted({i.rally_number for i in issues if i.rally_number is not None})
489         clean = n - len(flagged)
490         lines = [
491             f"# Delivery validation ? {source}",
492             "",
493             f"- rallies: **{n}**",
494             f"- issues: **{len(issues)}** (blocker {len(by_sev['blocker'])}, "
495             f"high {len(by_sev['high'])}, medium {len(by_sev['medium'])}, low
{len(by_sev['low'])})",
496             f"- structurally-clean rallies: **{clean}/{n}**"
497             + (f" ({100 * clean // n}%" if n else ""),
498             f"- gate: **{'PASS' if not by_sev['blocker'] else 'BLOCKED'}**",
499             "",
500         ]
501         if issues:
502             for sev in ("blocker", "high", "medium", "low"):
503                 grp = by_sev[sev]
504                 if not grp:
505                     continue
506                 lines.append(f"## {sev.upper()}")
507                 lines += [f"- ({i.code}) {i.message}" for i in grp]
508                 lines.append("")
509         else:
510             lines.append("No issues detected. ?\n")
511         if flagged:
512             lines.append(f"## Rallies to review: {flagged}")
513         return "\n".join(lines)
514
515
516     def write_review_csv(records: List[RallyRecord], issues: List[Issue], out_path: str) ->
None:
517         """Auto-generate the annotatable manual-verification sheet (blank `true_` cols)."""
518         import csv
519         flags: Dict[int, List[str]] = {}
520         for i in issues:
521             if i.rally_number is not None:
522                 flags.setdefault(i.rally_number, []).append(i.code)
523         os.makedirs(os.path.dirname(os.path.abspath(out_path)), exist_ok=True)
524         cols = ["rally_number", "start_mmss", "end_mmss", "dur_s", "shots_count",
525                 "ending_reason", "last_shot_outcome", "auto_flags",
526                 "VERIFIED(Y/N)", "true_start_mmss", "true_end_mmss",
527                 "true_ending_reason", "true_last_shot_outcome", "reviewer_comment"]
528         with open(out_path, "w", newline="", encoding="utf-8") as f:
529             w = csv.writer(f)
530             w.writerow(cols)
531             for r in sorted(records, key=lambda x: x.rally_number):
532                 w.writerow([

```

```

533         r.rally_number, _mmss(r.start_time), _mmss(r.end_time),
534         f"{r.duration:.1f}", r.shots_count if r.shots_count is not None else "",
535         r.ending_reason or "", r.last_shot_outcome or "",
536         ";".join(flags.get(r.rally_number, [])),
537         "", "", "", "", "", "", "",
538     ])
539

```

15. user (2026-06-10T06:00:41.761Z)

```

1     """Normalize a VLC rally-annotator CSV into a clean lean CSV that ``import-local``
2     ingests safely.
3
4     Two verified landmines in PARTIAL human annotation data (a video tagged only up to
5     some point, with mid-session corrections):
6
7     1. An unusable time aborts the WHOLE import (one bad row kills the entire video,
8     via the try/except around _parse_local_annotations in curate.py import-local):
9     a blank/missing time makes ``safe_float_required`` raise; a reversed
10    ``end <= start`` or a negative ``start`` makes the ``BadmintonRally`` pydantic
11    validator raise. -> we DROP such rows, mirroring the downstream indexer loader
12    which already skips ``not end > start``.
13    2. A fractional or duplicate ``rally_number`` (e.g. ``7.5`` inserted for a rally
14    the annotator missed) truncates via ``int(float())`` and COLLIDES on the
15    ``(video_id, rally_number)`` primary key; the DELETE-then-INSERT upsert then
16    silently destroys one rally. -> we RENUMBER 1..N by start_time, recording the
17    original id for traceability.
18
19    It also accepts mm:ss / HH:MM:SS times (converted to decimal seconds) and folds a
20    pipe-delimited ``ending_reason`` ('forced_error|other') to its primary token, so the
21    output is always the canonical lean schema::
22
23        rally_number,start_time,end_time,ending_reason,sport[,shots_count]
24
25    plus traceability columns (``original_rally_number``, ``notes``) that import-local
26    ignores. Standalone on purpose: it does NOT touch the DB schema or the
27    rally_number-ordered export. Run it, then feed the output to ``curate import-local``.
28
29    python -m src.cli.normalize_annotations --in "Match.rallies.csv" --out clean.csv
30    """
31    from __future__ import annotations
32
33    import argparse
34    import csv
35    import os
36    from dataclasses import dataclass, field
37    from typing import Any, Dict, List, Optional, Tuple
38
39    # Canonical lean output columns + traceability (import-local reads the first four/
40    # five; original_rally_number + notes are ignored by the importer but kept for humans).
41    OUT_FIELDS = ["rally_number", "start_time", "end_time", "ending_reason", "sport",
42                  "shots_count", "original_rally_number", "notes"]
43
44
45    @dataclass
46    class NormalizeReport:
47        rows_in: int = 0
48        rows_out: int = 0
49        dropped: List[Dict[str, Any]] = field(default_factory=list) # {orig_id, reason,
start, end}
50        renumbered: bool = False
51        fractional_ids: List[str] = field(default_factory=list)
52        duplicate_ids: List[str] = field(default_factory=list)
53        mmss_converted: int = 0
54

```



```

55     def summary(self) -> str:
56         lines = [f"rows in: {self.rows_in} -> rows out: {self.rows_out}"]
57         if self.dropped:
58             lines.append(f"dropped {len(self.dropped)} bad row(s):")
59             for d in self.dropped:
60                 lines.append(f"  - orig id {d['orig_id']!r}: {d['reason']} "
61                             f"(start={d['start']!r}, end={d['end']!r})")
62         if self.renumbered:
63             extra = []
64             if self.fractional_ids:
65                 extra.append(f"fractional ids {self.fractional_ids}")
66             if self.duplicate_ids:
67                 extra.append(f"duplicate ids {self.duplicate_ids}")
68             lines.append("renumbered rallies 1..N by start_time"
69                         + (f" ({'; '.join(extra)})" if extra else ""))
70         if self.mmss_converted:
71             lines.append(f"converted {self.mmss_converted} mm:ss timestamp(s) to
seconds")
72         if not self.dropped and not self.renumbered and not self.mmss_converted:
73             lines.append("clean input ? no changes beyond column normalization")
74         return "\n".join(lines)
75
76
77     def _to_seconds(raw: Any) -> Optional[Tuple[float, bool]]:
78         """(seconds, was_mmss) or None if blank/unparseable. Decimal seconds or mm:ss /
HH:MM:SS."""
79         if raw is None:
80             return None
81         s = str(raw).strip()
82         if not s:
83             return None
84         try:
85             if ":" in s:
86                 parts = [float(p) for p in s.split(":")]
87                 if len(parts) == 3:
88                     return parts[0] * 3600 + parts[1] * 60 + parts[2], True
89                 if len(parts) == 2:
90                     return parts[0] * 60 + parts[1], True
91                 return None
92             return float(s), False
93         except ValueError:
94             return None
95
96
97     def _primary_reason(raw: Any) -> Tuple[str, str]:
98         """Pipe-delimited 'forced_error|other' -> ('forced_error', 'other'). Blank -> ('',
'')."""
99         if raw is None:
100             return "", ""
101         s = str(raw).strip()
102         if not s:
103             return "", ""
104         parts = [p.strip() for p in s.split("|") if p.strip()]
105         if not parts:
106             return "", ""
107         return parts[0], "; ".join(parts[1:])
108
109
110     def _get(row: Dict[str, Any], *keys: str) -> Any:
111         """First non-empty value among the given (already-lowercased) keys."""
112         for k in keys:
113             v = row.get(k)
114             if v is not None and str(v).strip() != "":
115                 return v
116         return None

```

```

117
118
119     def normalize_rows(rows: List[Dict[str, Any]]) -> Tuple[List[Dict[str, Any]],
NormalizeReport]:
120         """Pure core: lowercase-keyed annotator rows -> clean lean rows + report."""
121         rep = NormalizeReport(rows_in=len(rows))
122         kept: List[Dict[str, Any]] = []
123
124         for row in rows:
125             orig_id = _get(row, "rally_number", "rally")
126             s_parsed = _to_seconds(_get(row, "start_time", "start_mmss", "start"))
127             e_parsed = _to_seconds(_get(row, "end_time", "end_mmss", "end"))
128             start = s_parsed[0] if s_parsed else None
129             end = e_parsed[0] if e_parsed else None
130
131             if start is None:
132                 rep.dropped.append({"orig_id": orig_id, "reason": "missing/invalid
start_time",
133                                     "start": _get(row, "start_time", "start_mmss"), "end":
_get(row, "end_time", "end_mmss")})
134                 continue
135             if start < 0:
136                 # BadmintonRally's pydantic validator rejects start<0 -> would abort the
137                 # whole import; drop it here like any other unusable row.
138                 rep.dropped.append({"orig_id": orig_id, "reason": "negative start_time",
139                                     "start": start, "end": end})
140                 continue
141             if end is None:
142                 rep.dropped.append({"orig_id": orig_id, "reason": "missing/invalid
end_time",
143                                     "start": start, "end": _get(row, "end_time",
"end_mmss")})
144                 continue
145             if end <= start:
146                 rep.dropped.append({"orig_id": orig_id, "reason": "end_time <= start_time",
147                                     "start": start, "end": end})
148                 continue
149             if (s_parsed and s_parsed[1]):
150                 rep.mmss_converted += 1
151             if (e_parsed and e_parsed[1]):
152                 rep.mmss_converted += 1
153
154             reason, reason_note = _primary_reason(_get(row, "ending_reason"))
155             note_in = _get(row, "notes", "note", "reviewer_comment")
156             notes = "; ".join([x for x in [reason_note, (str(note_in).strip() if note_in
else "")] if x])
157             kept.append({
158                 "orig_id": orig_id, "start": float(start), "end": float(end),
159                 "ending_reason": reason, "sport": _get(row, "sport"),
160                 "shots_count": _get(row, "shots_count", "shots"), "notes": notes,
161             })
162
163         kept.sort(key=lambda r: r["start"])
164
165         # Renumber 1..N by start order; detect whether ids were fractional / duplicate /
166         # not already a clean 1..N sequence (any of which means import-local would collide).
167         orig_seq = [str(r["orig_id"]) if r["orig_id"] is not None else "" for r in kept]
168         new_seq = [str(i) for i in range(1, len(kept) + 1)]
169         rep.renumbered = orig_seq != new_seq
170         seen: Dict[str, int] = {}
171         for oid in orig_seq:
172             seen[oid] = seen.get(oid, 0) + 1
173         rep.duplicate_ids = sorted({oid for oid, c in seen.items() if c > 1 and oid})
174         rep.fractional_ids = sorted({oid for oid in orig_seq if "." in oid})
175

```

```

176     out: List[Dict[str, Any]] = []
177     for i, r in enumerate(kept, start=1):
178         out.append({
179             "rally_number": i,
180             "start_time": f"{r['start']:.3f}",
181             "end_time": f"{r['end']:.3f}",
182             "ending_reason": r["ending_reason"] or "",
183             "sport": r["sport"] or "",
184             "shots_count": "" if r["shots_count"] is None else
str(r["shots_count"]).strip(),
185             "original_rally_number": "" if r["orig_id"] is None else str(r["orig_id"]),
186             "notes": r["notes"],
187         })
188     rep.rows_out = len(out)
189     return out, rep
190
191
192 def normalize_file(in_path: str, out_path: str) -> NormalizeReport:
193     with open(in_path, "r", newline="", encoding="utf-8-sig") as f:
194         reader = csv.DictReader(f)
195         reader.fieldnames = [(n or "").lower().strip() for n in (reader.fieldnames or
[[]])]
196         rows = [{(k or "").lower().strip(): v for k, v in r.items()} for r in reader]
197     out, rep = normalize_rows(rows)
198     os.makedirs(os.path.dirname(os.path.abspath(out_path)) or ".", exist_ok=True)
199     with open(out_path, "w", newline="", encoding="utf-8") as f:
200         w = csv.DictWriter(f, fieldnames=OUT_FIELDS)
201         w.writeheader()
202         w.writerows(out)
203     return rep
204
205
206 def main() -> None:
207     ap = argparse.ArgumentParser(description="Normalize a rally-annotator CSV for safe
import-local.")
208     ap.add_argument("--in", dest="in_path", required=True, help="rally-annotator
.rallies.csv")
209     ap.add_argument("--out", dest="out_path", required=True, help="clean lean CSV to
write")
210     args = ap.parse_args()
211     rep = normalize_file(args.in_path, args.out_path)
212     print("=== normalize-annotations ===")
213     print(rep.summary())
214     print(f"wrote -> {args.out_path}")
215
216
217 if __name__ == "__main__":
218     main()
219

```

16. user (2026-06-10T06:00:42.923Z)

```

1     """Shared parsing helpers used by the CLI importer and dataset parsers.
2
3     Previously these were duplicated as private ``_safe_int`` / ``_safe_float`` /
4     ``_parse_time_string`` in both ``src/cli/curate.py`` and the CoachAI parser.
5     """
6     from typing import Any, Optional
7
8
9     def safe_int(value: Any, default: Optional[int] = 0) -> Optional[int]:
10         """Parse an int from a possibly-empty/float/None value; return default on
failure."""
11         if value is None:
12             return default

```

```

13         s = str(value).strip()
14         if not s:
15             return default
16         try:
17             return int(float(s))
18         except (TypeError, ValueError):
19             return default
20
21
22     def safe_float_required(value: Any) -> float:
23         """Parse a float; raise ValueError if missing/empty (caller treats it as
required)."""
24         if value is None:
25             raise ValueError("missing required numeric value")
26         s = str(value).strip()
27         if not s:
28             raise ValueError("missing required numeric value")
29         return float(s)
30
31
32     def parse_time_string(time_str: Any) -> float:
33         """Convert 'HH:MM:SS' / 'MM:SS' / float-strings to float seconds; 0.0 if
blank/invalid."""
34         if not time_str:
35             return 0.0
36         s = str(time_str).strip()
37         if ":" in s:
38             parts = s.split(":")
39             try:
40                 if len(parts) == 3:
41                     return float(parts[0]) * 3600 + float(parts[1]) * 60 + float(parts[2])
42                 if len(parts) == 2:
43                     return float(parts[0]) * 60 + float(parts[1])
44             except ValueError:
45                 return 0.0
46         try:
47             return float(s)
48         except ValueError:
49             return 0.0
50

```

17. user (2026-06-10T06:00:43.024Z)

```

1     from typing import List, Tuple, Union
2     from src.core.config import load_config
3     from src.core.schemas.base import VideoMetadata, LicenseType
4     from src.core.schemas.badminton import BadmintonRally
5     from src.core.schemas.tennis import TennisPoint
6     from src.core.schemas.cricket import CricketDelivery
7     from src.core.schemas.pickleball import PickleballRally
8     from src.core.schemas.tabletennis import TableTennisRally
9
10    class DatasetValidator:
11        def __init__(self, config_path: str = "config/config.yaml"):
12            self.config_path = config_path
13            self.commercial_licenses = self._load_licenses_whitelist()
14
15        def _load_licenses_whitelist(self) -> List[str]:
16            # load_config returns {} on a missing/unreadable file, so the default list
17            # below is the single fallback (keep in sync with config.yaml).
18            config = load_config(self.config_path)
19            return config.get(
20                'commercial_licenses',
21                ["MIT", "Apache-2.0", "BSD", "CC0", "CC-BY", "Public-Domain",
"Proprietary"],

```

```

22         )
23
24     def check_license_commercial(self, license_type: Union[LicenseType, str]) -> bool:
25         """Determines if a license type is safe for commercial use based on the
configuration whitelist."""
26         val = license_type.value if isinstance(license_type, LicenseType) else
license_type
27         return val in self.commercial_licenses
28
29     def validate_badminton_rallies(self, rallies: List[BadmintonRally]) -> Tuple[bool,
List[str]]:
30         """
31         Validates badminton rallies:
32         - Sorted chronologically.
33         - Start time < End time (handled by Pydantic, but double checked).
34         - No temporal overlap between sequential rallies.
35         """
36         errors = []
37         if not rallies:
38             return True, errors
39
40         # Sort by rally number to check progression
41         sorted_rallies = sorted(rallies, key=lambda r: r.rally_number)
42
43         for i in range(len(sorted_rallies)):
44             r = sorted_rallies[i]
45             if r.start_time >= r.end_time:
46                 errors.append(f"Rally {r.rally_number}: start_time ({r.start_time}) >=
end_time ({r.end_time})")
47
48             if i > 0:
49                 prev = sorted_rallies[i - 1]
50                 if r.start_time < prev.end_time:
51                     errors.append(
52                         f"Rally overlap detected between Rally {prev.rally_number} "
53                         f"(ends at {prev.end_time}s) and Rally {r.rally_number} "
54                         f"(starts at {r.start_time}s)"
55                     )
56
57         return len(errors) == 0, errors
58
59     def validate_tennis_points(self, points: List[TennisPoint]) -> Tuple[bool,
List[str]]:
60         errors = []
61         if not points:
62             return True, errors
63
64         sorted_pts = sorted(points, key=lambda p: p.point_number)
65         for i in range(len(sorted_pts)):
66             pt = sorted_pts[i]
67             if pt.start_time >= pt.end_time:
68                 errors.append(f"Point {pt.point_number}: start ({pt.start_time}) >= end
({pt.end_time})")
69
70             if i > 0:
71                 prev = sorted_pts[i - 1]
72                 if pt.start_time < prev.end_time:
73                     errors.append(
74                         f"Point overlap detected: Point {prev.point_number} (ends
{prev.end_time}s) "
75                         f"and Point {pt.point_number} (starts {pt.start_time}s)"
76                     )
77         return len(errors) == 0, errors
78
79     def validate_cricket_deliveries(self, deliveries: List[CricketDelivery]) ->

```

```

Tuple[bool, List[str]]:
80         errors = []
81         if not deliveries:
82             return True, errors
83
84         # Sort by over, then ball within over
85         sorted_delivs = sorted(deliveries, key=lambda d: (d.over, d.ball))
86         for i in range(len(sorted_delivs)):
87             d = sorted_delivs[i]
88             if d.start_time >= d.end_time:
89                 errors.append(f"Over {d.over} Ball {d.ball}: start ({d.start_time}) >=
end ({d.end_time})")
90
91             if i > 0:
92                 prev = sorted_delivs[i - 1]
93                 if d.start_time < prev.end_time:
94                     errors.append(
95                         f"Delivery overlap: Over {prev.over} Ball {prev.ball} (ends
{prev.end_time}s) "
96                         f"and Over {d.over} Ball {d.ball} (starts {d.start_time}s)"
97                     )
98             return len(errors) == 0, errors
99
100     def validate_pickleball_rallies(self, rallies: List[PickleballRally]) -> Tuple[bool,
List[str]]:
101         """Validates pickleball rallies: positive durations and no temporal overlap
102         between chronologically-ordered rallies (same contract as badminton)."""
103         errors = []
104         if not rallies:
105             return True, errors
106
107         sorted_rallies = sorted(rallies, key=lambda r: r.rally_number)
108         for i in range(len(sorted_rallies)):
109             r = sorted_rallies[i]
110             if r.start_time >= r.end_time:
111                 errors.append(f"Rally {r.rally_number}: start_time ({r.start_time}) >=
end_time ({r.end_time})")
112
113             if i > 0:
114                 prev = sorted_rallies[i - 1]
115                 if r.start_time < prev.end_time:
116                     errors.append(
117                         f"Rally overlap detected between Rally {prev.rally_number} "
118                         f"(ends at {prev.end_time}s) and Rally {r.rally_number} "
119                         f"(starts at {r.start_time}s)"
120                     )
121             return len(errors) == 0, errors
122
123     def validate_tabletennis_rallies(self, rallies: List[TableTennisRally]) ->
Tuple[bool, List[str]]:
124         """Validates table-tennis rallies: positive durations and no temporal overlap
125         between chronologically-ordered rallies (same contract as badminton)."""
126         errors = []
127         if not rallies:
128             return True, errors
129
130         sorted_rallies = sorted(rallies, key=lambda r: r.rally_number)
131         for i in range(len(sorted_rallies)):
132             r = sorted_rallies[i]
133             if r.start_time >= r.end_time:
134                 errors.append(f"Rally {r.rally_number}: start_time ({r.start_time}) >=
end_time ({r.end_time})")
135
136             if i > 0:
137                 prev = sorted_rallies[i - 1]

```

```

138         if r.start_time < prev.end_time:
139             errors.append(
140                 f"Rally overlap detected between Rally {prev.rally_number} "
141                 f"(ends at {prev.end_time}s) and Rally {r.rally_number} "
142                 f"(starts at {r.start_time}s)"
143             )
144     return len(errors) == 0, errors
145

```

18. user (2026-06-10T06:00:49.776Z)

```

1     from pydantic import BaseModel, Field
2     from typing import Optional
3     from enum import Enum
4
5     class SportType(str, Enum):
6         BADMINTON = "badminton"
7         TENNIS = "tennis"
8         TABLE_TENNIS = "tabletennis"
9         PICKLEBALL = "pickleball"
10        CRICKET = "cricket"
11
12    class LicenseType(str, Enum):
13        MIT = "MIT"
14        APACHE_2 = "Apache-2.0"
15        BSD = "BSD"
16        CC0 = "CC0"
17        CC_BY = "CC-BY"
18        CC_BY_NC = "CC-BY-NC"
19        CC_BY_NC_SA = "CC-BY-NC-SA"
20        PUBLIC_DOMAIN = "Public-Domain"
21        PROPRIETARY = "Proprietary"
22        UNKNOWN = "Unknown"
23
24    class CurationStatus(str, Enum):
25        STAGING = "staging"
26        CURATED = "curated"
27        REJECTED = "rejected"
28
29    class VideoMetadata(BaseModel):
30        video_id: str = Field(..., description="Unique alphanumeric identifier for the
video")
31        sport: SportType
32        video_url: Optional[str] = Field(None, description="Remote URL of the video (if
available)")
33        local_path: Optional[str] = Field(None, description="Local path to video file (if
downloaded or custom)")
34        license_type: LicenseType
35        license_url: Optional[str] = None
36        is_commercial: bool = Field(False, description="Flag indicating if the license
allows commercial use")
37        status: CurationStatus = CurationStatus.STAGING
38        match_name: str = Field(..., description="Name of the match/event")
39        fps: Optional[float] = Field(None, description="Frames per second")
40        resolution: Optional[str] = Field(None, description="Resolution string (e.g. 1080p,
720p)")
41

```

19. user (2026-06-10T06:00:50.739Z)

```

1     from pydantic import BaseModel, Field, model_validator
2     from typing import Optional
3
4     class BadmintonRally(BaseModel):
5         video_id: str

```

```

6         rally_number: int = Field(..., description="Chronological rally index in the match")
7         start_time: float = Field(..., description="Start of rally in seconds")
8         end_time: float = Field(..., description="End of rally in seconds")
9         score_before: Optional[str] = Field(None, description="Score before the rally (e.g.
'0-0')")
10        score_after: Optional[str] = Field(None, description="Score after the rally (e.g.
'1-0')")
11        shots_count: Optional[int] = Field(None, description="Number of strokes in the
rally")
12        ending_reason: Optional[str] = Field(None, description="Why the rally ended:
'winner' | 'forced_error' | 'unforced_error' | 'service_fault' | 'let' | 'other'")
13        last_shot_outcome: Optional[str] = Field(None, description="Where the last shot
landed: 'in' | 'net' | 'out' | 'n_a' (n_a for service_fault/let)")
14
15        @model_validator(mode='after')
16        def validate_timestamps(self) -> 'BadmintonRally':
17            if self.start_time < 0:
18                raise ValueError("start_time must be non-negative")
19            if self.end_time <= self.start_time:
20                raise ValueError("end_time must be strictly greater than start_time")
21            return self
22

```

20. user (2026-06-10T06:00:51.210Z)

```

1         """Central configuration loader.
2
3         A single place to read ``config/config.yaml`` so modules don't each re-open the
4         file with slightly different fallbacks. Returns ``{}`` when the file is missing or
5         unreadable (callers apply their own defaults), matching the prior per-module
6         behaviour but without the duplicated boilerplate and broad excepts.
7         """
8         import os
9         from typing import Any, Dict
10
11         import yaml
12
13         DEFAULT_CONFIG_PATH = "config/config.yaml"
14         DEFAULT_DB_PATH = "data/sports_metadata.db"
15
16
17         def load_config(config_path: str = DEFAULT_CONFIG_PATH) -> Dict[str, Any]:
18             """Load a YAML config into a dict; return {} if missing/unreadable/empty."""
19             if not os.path.exists(config_path):
20                 return {}
21             try:
22                 with open(config_path, "r", encoding="utf-8") as f:
23                     return yaml.safe_load(f) or {}
24             except (OSError, yaml.YAMLError):
25                 return {}
26
27
28         def get_database_path(config: Dict[str, Any]) -> str:
29             """Resolve the SQLite database path from a loaded config, with a default."""
30             return config.get("database_path", DEFAULT_DB_PATH)
31

```

21. user (2026-06-10T06:00:51.398Z)

```

1         """Adapt collector data into an obsreport envelope.
2
3         Reuses the existing rally_cvat report primitives (``Issue``, records, severities)
4         rather than re-deriving anything ? this module only *maps* them onto the shared
5         report shape. Imported only after ``src.reporting`` confirms obsreport is present,
6         so it may import obsreport freely.

```



```

7      """
8      from __future__ import annotations
9
10     from pathlib import Path
11     from typing import Any, List, Optional
12
13     import obsreport
14
15     _SPEC_PATH = Path(__file__).resolve().parent / "spec.yaml"
16
17     # Delivery issue severity -> report flag severity.
18     _SEVERITY = {"blocker": "error", "high": "warn", "medium": "warn", "low": "info"}
19     # Every issue code is documented in this FAQ table.
20     _ISSUE_FAQ = "docs/INGESTION_PIPELINE.md#what-does-validation-check"
21
22
23     def load_spec() -> "obsreport.ReportSpec":
24         return obsreport.load_spec(_SPEC_PATH)
25
26
27     # Heuristic checks that flag a region to REVIEW ? not hard failures. We frame these
28     # gently for the vendor so a "possible" finding isn't read as a certainty/blocker.
29     _REVIEW_ONLY = {"possible_missed_rally", "shot_density"}
30
31
32     def issue_to_flag(issue: Any, record: Any = None, fps: Any = None) -> "obsreport.Flag":
33         """Map a rally_cvat.Issue onto a report Flag. Actionable issues (blocker/high/
34         medium) also carry a customer_message so the delivery summary shows the vendor
35         exactly what to fix, in plain language. When the rally's source record + fps are
36         known, frame-level evidence is attached so a CVAT (frame-based) reviewer can act
37         without converting seconds to frames by hand."""
38         sev = _SEVERITY.get(issue.severity, "info")
39         customer = issue.message if sev in ("error", "warn") else None
40         if customer and issue.code in _REVIEW_ONLY:
41             customer += (" (Review only ? a heuristic, not a blocker: check the video and "
42                         "act only if confirmed.)")
43         evidence: dict = {}
44         if issue.rally_number is not None:
45             evidence["rally_number"] = issue.rally_number
46         if record is not None and getattr(record, "frame_min", -1) is not None and
47         getattr(record, "frame_min", -1) >= 0:
48             evidence["rally_frames"] = f"{record.frame_min}-{record.frame_max}"
49         if fps:
50             evidence["fps"] = fps
51         return obsreport.Flag(
52             code=issue.code,
53             severity=sev,
54             message=issue.message,
55             stage="validation",
56             faq_ref=_ISSUE_FAQ,
57             customer_message=customer,
58             evidence=evidence,
59         )
60
61     def _recorder(source: str, video_path: Optional[str], spec) -> "obsreport.RunRecorder":
62         return obsreport.RunRecorder(
63             "sports-data-collector",
64             video_id=source,
65             video_id_kind="human_id",
66             video_path=video_path,
67             spec=spec if spec is not None else load_spec(),
68         )
69
70

```

```

71 def build_delivery_recorder(
72     *,
73     records: List[Any],
74     issues: List[Any],
75     fps: Optional[float],
76     source: str,
77     video_path: Optional[str] = None,
78     video_duration_s: Optional[float] = None,
79     source_format: str = "cvat",
80     step: Optional[int] = None,
81     spec=None,
82 ) -> "obsreport.RunRecorder":
83     """Envelope for `validate-delivery` / `convert-cvat`: parse + validation stages,
84     one flag per delivery issue, a rally summary."""
85     rec = _recorder(source, video_path, spec)
86     if fps:
87         rec.set_video_meta(fps=fps, fps_source="declared")
88     if video_duration_s:
89         rec.set_video_meta(duration_s=round(video_duration_s, 1))
90
91     with rec.stage("parse") as st:
92         st.add_metric("rallies_parsed", len(records))
93         st.details.update({"source_format": source_format, "fps": fps, "sampling_step":
step})
94
95     blockers = [i for i in issues if i.severity == "blocker"]
96     with rec.stage("validation") as st:
97         for sev in ("blocker", "high", "medium", "low"):
98             st.add_metric(f"issues_{sev}", sum(1 for i in issues if i.severity == sev))
99         st.add_metric("n_blockers", len(blockers))
100        st.details["gate"] = "BLOCKED" if blockers else "PASS"
101        st.status = "error" if blockers else ("warn" if issues else "ok")
102
103    records_by_rally = {getattr(r, "rally_number", None): r for r in records}
104    for issue in issues:
105        record = records_by_rally.get(issue.rally_number)
106        rec.envelope.flags.append(issue_to_flag(issue, record=record, fps=fps))
107
108    _rally_summary(rec, records, video_duration_s)
109    rec.finalize()
110    return rec
111
112
113 def build_import_recorder(
114     *,
115     video_meta_obj: Any,
116     annotations: List[Any],
117     is_valid: bool,
118     errors: List[str],
119     count: int,
120     sport: str,
121     source_format: str,
122     spec=None,
123 ) -> "obsreport.RunRecorder":
124     """Envelope for `import-local`: parse + validation + registration stages."""
125     rec = _recorder(video_meta_obj.video_id, video_meta_obj.local_path, spec)
126     fps = getattr(video_meta_obj, "fps", None)
127     rec.set_video_meta(
128         fps=fps,
129         fps_source="declared" if fps else "unknown",
130         resolution=getattr(video_meta_obj, "resolution", None),
131     )
132
133     with rec.stage("parse") as st:
134         st.add_metric("records_parsed", len(annotations))

```

```

135         st.details["source_format"] = source_format
136
137     with rec.stage("validation") as st:
138         st.status = "ok" if is_valid else "error"
139         st.add_metric("valid", is_valid)
140         for e in errors or []:
141             st.messages.append(str(e))
142
143     with rec.stage("registration") as st:
144         st.add_metric("segments_registered", count)
145         st.details.update({
146             "sport": sport,
147             "license": getattr(getattr(video_meta_obj, "license_type", None), "value",
148                               str(getattr(video_meta_obj, "license_type", ""))),
149             "is_commercial": bool(getattr(video_meta_obj, "is_commercial", False)),
150             "status": getattr(getattr(video_meta_obj, "status", None), "value",
151                               str(getattr(video_meta_obj, "status", ""))),
152         })
153
154     _rally_summary(rec, annotations, None)
155     rec.finalize()
156     return rec
157
158
159     def _rally_summary(rec: "obsreport.RunRecorder", records: List[Any], video_duration_s:
Optional[float]) -> None:
160         total = 0.0
161         for r in records:
162             try:
163                 total += float(r.end_time) - float(r.start_time)
164             except (AttributeError, TypeError, ValueError):
165                 pass
166         coverage = round(100.0 * total / video_duration_s, 1) if video_duration_s else None
167         notes = [] if records else ["No rallies were found in this delivery."]
168         rec.set_highlights(segment_count=len(records), total_highlight_s=round(total, 1),
169                           coverage_pct=coverage, notes=notes)
170

```

22. assistant (2026-06-10T06:00:59.203Z)

Now let me run the test suite and read the remaining pieces in parallel.

23. user (2026-06-10T06:01:04.807Z)

```

..... [ 58%]
..... [100%]
124 passed in 2.88s

```

24. user (2026-06-10T06:01:04.833Z)

```

1     # Roora Annotation Brief ? Rally Golden Set
2
3     **Project:** racket-sport rally segmentation golden / regression dataset
4     **Sports in scope:** Badminton, Pickleball, Table Tennis (each in **singles _and_
doubles**)
5     **Companion doc:** [ANNOTATION_GUIDE.md](./ANNOTATION_GUIDE.md)
6
7     > This brief is the vendor-facing spec. The annotation deliverable is a per-rally CSV
that
8     > imports directly via `python -m src.cli.curate import-local`.
9
10    ## 0. What we need in one sentence
11    For each provided match video, produce **one CSV file** listing **every rally** with its
12    start time, end time, number of shots, and how it ended ? **excluding all dead time**.

```

```

13
14 ## 1. Deliverable format (canonical)
15 One CSV per video (UTF-8), named exactly `<video_id>.csv` (we supply the `video_id`).
16 One row per rally, chronological, header row required.
17
18 | column | meaning |
19 |---|---|
20 | `rally_number` | integer, 1-based, chronological. No gaps, no repeats. |
21 | `start_mmss` | rater-entered time as `m:ss.mmm` (e.g. `1:12.40`) ? serve contact. |
22 | `end_mmss` | rater-entered time as `m:ss.mmm` ? ball/shuttle dead. |
23 | `start_time` | decimal seconds (e.g. `72.40`). Auto-filled from `start_mmss`.
24 | |
25 | `end_time` | decimal seconds. Auto-filled from `end_mmss`. |
26 | `shots_count` | integer. Total racket/paddle/bat contacts (serve = shot 1). |
27 | `ending_reason` | one of: `winner` \| `forced_error` \| `unforced_error` \|
28 | `service_fault` \| `let` \| `other` (why it ended) |
29 | `last_shot_outcome` | one of: `in` \| `net` \| `out` \| `n_a` (where the last shot
30 | landed; `n_a for service_fault/let) |
31 | `score_before` | optional, `"A-B". Blank if not confidently readable. |
32 | `score_after` | optional, `"A-B". |
33 | `notes` | optional free text. Required when `ending_reason = other`. |
34
35 Times (mm:ss ? decimal seconds). Raters type only `start_mmss`/`end_mmss`. The
36 template
37 Google Sheet fills `start_time`/`end_time` with this formula (shown for `start_time`,
38 row 2):
39
40 ```
41 =IF(B2="", "", IFERROR(VALUE(LEFT(B2,FIND(":",B2)-1))*60 +
42 VALUE(MID(B2,FIND(":",B2)+1,20)), VALUE(B2)))
43 ```
44
45 On Download as CSV, the computed decimal seconds are exported. (If your tool already
46 shows decimal seconds, type those into `start_time`/`end_time` and leave the mm:ss
47 columns blank.)
48
49 Example (badminton):
50
51 ```csv
52 rally_number,start_mmss,end_mmss,start_time,end_time,shots_count,ending_reason,last_shot
53 _outcome,score_before,score_after,notes
54 1,0:12.50,0:25.00,12.50,25.00,9,winner,in,0-0,1-0,
55 2,0:45.00,0:47.30,45.00,47.30,1,service_fault,n_a,1-0,1-1,serve into net
56 3,1:01.20,1:18.85,61.20,78.85,14,unforced_error,out,1-1,2-1,final clear sailed long (no
57 pressure)
58 ```
59
60 Critical format rules
61
62 - `start_time`/`end_time` that reach us must be decimal seconds (the formula
63 guarantees this).
64
65 - Times are relative to the start of the file we provide. Do not trim, clip, or
66 re-encode the video ? our pipeline identifies a video by its file checksum.
67
68 - `end_time` > `start_time` for every row.
69
70 - No overlaps: `end_time[N] <= start_time[N+1]`.
71
72 - Everything between one rally's `end_time` and the next rally's `start_time` is dead
73 time and gets no row.
74
75
76 ## 2. The three non-negotiable rules (full detail + pictures in the Guide)
77
78 1. Ignore all dead time. Only label active play.
79
80 2. Include every rally ? exclude none. Every served point, including service
81 faults,
82
83 1?2 shot rallies, and lets/replays. A missed rally is the worst error in this
84 project.
85
86 Exception: warm-up / knock-up is NOT a rally ? pre-match and between-games warm-
87 up looks
88
89 like rallying but isn't scored play, so don't label it. Tell it apart by: no real

```

```

serve,
64     cooperative (not competitive) hitting, the score not progressing, often several
shuttles/balls
65     in use (full how-to in the Guide, Part 2). When unsure if the match has started, flag
it.
66     3. **A shot = one contact** with the racket/paddle/bat. Serve = shot 1. Bounces and net-
clips are not shots.
67
68     ## 3. Scope & phasing
69     **Phase A ? Pilot (this engagement):** 3 videos (1 badminton + 1 pickleball + 1 table
tennis),
70     each `< 3 min` with ~10?12 rallies. This is the **paid calibration round**; we review
jointly and
71     lock the Guide before scale-up.
72
73     **Phase B ? Scale (after a successful pilot):** ~12?15 videos **per sport** (~36?45
total), each
74     ~30 min. Same format, rules, and template.
75
76     > **Doubles are included** across the program. Rally rules are identical to singles; the
77     > badminton- and pickleball-specific doubles notes are in the Guide.
78
79     ## 4. Quality bar (acceptance criteria)
80     Automated validation (no overlaps, positive durations, schema-valid) runs on every file,
then a
81     ~20% human spot-check. A file is returned for rework (in scope) if it misses:
82     - **Completeness:** every rally present; zero missed rallies; zero dead-time rows.
83     - **Boundaries:** `start_time`/`end_time` within **±0.3 s** of true serve-contact /
landing.
84     - **shots_count:** exact for rallies ? 15 shots; ±1 allowed for longer/faster rallies.
85     - **ending_reason:** one of the six allowed values, per the Guide's decision tree.
86     - **Format:** decimal seconds present, valid columns, non-overlapping, chronological.
87
88     ## 5. Per-sport quick reference (full version + examples in the Guide)
89     **Badminton (singles & doubles)** ? start: racket?shuttle serve contact · end: shuttle
floors / fault · shot: each racket?shuttle contact.
90     **Pickleball (singles & doubles)** ? start: paddle?ball serve contact · end: double
bounce / out / net / fault · shot: each paddle?ball contact (floor bounces are **not** shots;
two-bounce rule is normal play).
91     **Table tennis** ? start: racket?ball serve strike (not the toss) · end: failed legal
return · shot: each racket?ball contact (table bounces are **not** shots). Net-cord serve =
`let`.
92
93     ## 6. ending_reason values (decision tree in the Guide)
94     This vocabulary attributes **why** the point ended; it is shared across all
95     net-separated sports (badminton, pickleball, table tennis, tennis). Record **where**
96     the last shot physically landed in the separate **`last_shot_outcome`** column
97     (`in` \| `net` \| `out` \| `n_a`) ? e.g. a forced error that flew long is
98     `ending_reason=forced_error, last_shot_outcome=out`.
99
100    | value | meaning |
101    |---|---|
102    | `winner` | last shot landed **in** and the opponent made no legal contact (a clean
winner). |
103    | `forced_error` | the losing player erred (netted / out / missed) **under pressure** ?
stretched, wrong-footed, or rushed by the opponent's strong shot. |
104    | `unforced_error` | the losing player erred on a **routine, unpressured** ball they had
time and position to make. |
105    | `service_fault` | the serve itself failed / was illegal. |
106    | `let` | the point was replayed. |
107    | `other` | anything else / cannot tell. **Must** add a note. |
108
109    ## 7. Recommended capture method
110    Use any tool that shows current time (mm:ss.mmm or seconds, or frame + known fps) and
can

```

```

111     **step one frame** forward/back: a frame-steppable player (VLC frame-step) or a timeline
tool
112     (Label Studio / CVAT). The deliverable is the CSV in the schema above.
113     `seconds = frame_number / fps` (we provide fps per video).
114
115     ## 8. Logistics & contacts
116     - **TODO(avidullu): paste the Google Drive folder link** where source videos are shared
with Roora.
117     - Per-video `fps` + `video_id`: provided in the intake manifest ? template
[`INTAKE_MANIFEST_TEMPLATE.csv`](./INTAKE_MANIFEST_TEMPLATE.csv). **TODO(avidullu): fill the
manifest (one row per pilot video) and share its link.**
118     - **TODO(avidullu): paste the Drive folder link** where Roora uploads completed
`<video_id>.csv` files.
119     - Blank annotation template Sheet (mm:ss ? seconds formula from $1 pre-filled):
**TODO(avidullu): create it and paste the link.**
120     - **TODO(avidullu): add the escalation channel** (email / WhatsApp group) for edge-case
questions.
121     - NDA before any footage is shared ? template [`NDA_TEMPLATE.md`](./NDA_TEMPLATE.md).
**TODO(avidullu): have counsel review, fill the blanks, and have Roora sign.**
122     - Roora contact on file: `info@rooragrp.com` (HSR Layout, Bangalore).
123
124     ## 9. Pricing request (please quote)
125     1. Price **per rally** and **per minute** of footage.
126     2. **Pilot** price (3 short videos, ~30?36 rallies) as a paid calibration round.
127     3. Indicative **scale** price for ~36?45 longer videos.
128     4. Turnaround for the pilot and throughput (rallies/day) at scale.
129     5. QA process and rework policy.
130     6. Whether annotators have racket-sport familiarity (helpful, not required).
131     7. Confirmation you can deliver the CSV schema above.
132

```

25. user (2026-06-10T06:01:15.438Z)

```

backend\eval\__init__.py
backend\eval\annotations.py
backend\eval\classifier.py
backend\eval\metrics.py
backend\eval\training.py
backend\eval\calibration.py
backend\eval\gemini_refine.py
backend\eval\calibrate_local.py
backend\eval\calibrate_wasb.py
backend\eval\distill_local.py
backend\eval\export_wasb_segments.py
backend\eval\harness.py
backend\eval\regression.py
backend\eval\batch.py
backend\eval\eval_partial_labels.py
backend\eval\rally_seq_proto.py

```

26. user (2026-06-10T06:01:16.090Z)

```
run_phase3_eval.py
```

27. user (2026-06-10T06:01:22.224Z)

```

1     """Batch evaluation across a collector export manifest.
2
3     Phase 3 of the collector->indexer bridge: walk the `manifest.json` emitted by
4     sports-data-collector's `curate export`, (optionally) process each video through
5     a segmenter so detections land in the indexer DB, then score every video's
6     detected rallies against its ground-truth CSV and aggregate the results.
7
8     This module holds the *pure* pieces (path resolution, stage selection, metric

```

```

9      aggregation) so they can be unit-tested without touching video files, the GPU,
10     or any API. The orchestration/side-effects live in `run_phase3_eval.py`.
11     """
12
13     import os
14     from typing import Dict, List, Optional
15
16
17     # Segmenters that only ever populate the `final` play_segments table (no raw
18     # candidate stage), so scoring `candidates` for them would be meaninglessly empty.
19     _FINAL_ONLY_SEGMENTERS = {"motion", "gemini"}
20
21
22     def resolve_video_path(local_path: Optional[str], videos_root: str) -> Optional[str]:
23         """Resolve a manifest `local_path` to an absolute path.
24
25         The collector stores paths like ``./data/videos\\shuttle1_1.mp4`` relative
26         to its own repo root and with Windows separators. Normalise separators, strip
27         a leading ``./``, and resolve relative paths against `videos_root` (the
28         collector root). Returns None for a missing/empty path.
29         """
30         if not local_path:
31             return None
32         p = local_path.replace("\\", "/")
33         if p.startswith("./"):
34             p = p[2:]
35         # Cross-platform: treat Windows drive-letter absolute paths (C:/...) as absolute
36         # even when running on Linux (os.path.isabs("C:/...") == False on Unix).
37         # This keeps collector manifest paths working in tests and mixed environments.
38         if (len(p) > 1 and p[1] == ":" and p[0].isalpha()) or os.path.isabs(p):
39             return os.path.normpath(p)
40         return os.path.normpath(os.path.join(videos_root, p))
41
42
43     def default_stages_for(segmenter: str, requested: str = "auto") -> List[str]:
44         """Pick which prediction stages to score.
45
46         ``auto`` scores both stages for two-stage detectors (yolo_hybrid/tracknet_*)
47         but only ``final`` for segmenters that don't emit candidates. An explicit
48         request ('candidates' | 'final' | 'both') overrides the heuristic.
49         """
50         if requested == "both":
51             return ["candidates", "final"]
52         if requested in ("candidates", "final"):
53             return [requested]
54         # auto
55         if segmenter in _FINAL_ONLY_SEGMENTERS:
56             return ["final"]
57         return ["candidates", "final"]
58
59
60     def _safe_div(num: float, den: float) -> float:
61         return num / den if den else 0.0
62
63
64     def aggregate(per_video: List[dict], stages: List[str]) -> Dict[str, dict]:
65         """Aggregate per-video stage scores into corpus-level metrics.
66
67         `per_video` is a list of dicts shaped like `report_to_dict` output (each has
68         `stages[stage]["score"]` with tp/fp/fn/mean_iou/...). For each stage we
69         compute **micro** precision/recall/F1 (pooling TP/FP/FN across all videos) and
70         a TP-weighted mean IoU, plus the macro mean of per-video F1.
71         """
72         out: Dict[str, dict] = {}
73         for stage in stages:

```

```

74         tp = fp = fn = 0
75         iou_weighted = 0.0
76         f1s: List[float] = []
77         n_videos = 0
78         for rep in per_video:
79             sr = rep.get("stages", {}).get(stage)
80             if not sr:
81                 continue
82             s = sr["score"]
83             tp += s["true_positives"]
84             fp += s["false_positives"]
85             fn += s["false_negatives"]
86             iou_weighted += s["mean_iou"] * s["true_positives"]
87             f1s.append(s["f1"])
88             n_videos += 1
89         precision = _safe_div(tp, tp + fp)
90         recall = _safe_div(tp, tp + fn)
91         out[stage] = {
92             "videos": n_videos,
93             "tp": tp, "fp": fp, "fn": fn,
94             "micro_precision": precision,
95             "micro_recall": recall,
96             "micro_f1": _safe_div(2 * precision * recall, precision + recall),
97             "macro_f1": _safe_div(sum(f1s), len(f1s)),
98             "mean_iou": _safe_div(iou_weighted, tp),
99         }
100     return out
101
102
103 def format_aggregate(agg: Dict[str, dict]) -> str:
104     """Human-readable corpus-level summary table."""
105     lines = [
106         "=" * 78,
107         " CORPUS AGGREGATE",
108         "=" * 78,
109         f" {'Stage':<12} {'Vids':>4} {'TP':>5} {'FP':>5} {'FN':>5} "
110         f"{'mPrec':>6} {'mRec':>6} {'microF1':>8} {'macroF1':>8} {'mIoU':>6}",
111         "-" * 78,
112     ]
113     for stage, a in agg.items():
114         lines.append(
115             f" {stage:<12} {a['videos']:>4} {a['tp']:>5} {a['fp']:>5} {a['fn']:>5} "
116             f"{a['micro_precision']:>6.2f} {a['micro_recall']:>6.2f} "
117             f"{a['micro_f1']:>8.2f} {a['macro_f1']:>8.2f} {a['mean_iou']:>6.2f}"
118         )
119     lines.append("=" * 78)
120     return "\n".join(lines)
121

```

28. user (2026-06-10T06:01:27.394Z)

```

1     """Phase 3 driver: batch process + evaluate a collector export against ground truth.
2
3     Walks the `manifest.json` produced by sports-data-collector's `curate export`,
4     processes each video through a segmenter (so detections land in the indexer DB),
5     scores each against its ground-truth rally CSV, and prints per-video plus
6     corpus-aggregate precision/recall/F1, mean IoU, and boundary error.
7
8     Examples
9     -----
10    Free CPU baseline on a single video (validate the loop):
11        python run_phase3_eval.py --manifest ../sports-data-
12        collector/data/eval_export/manifest.json \
13        --segmenter motion --only shuttleset_1

```



```

14 Full corpus with the two-stage detector (needs GEMINI_API_KEY; costs API calls):
15     python run_phase3_eval.py --manifest ../manifest.json --segmenter yolo_hybrid \
16         --json output/phase3_report.json
17
18 Re-score only (videos already processed), no re-segmentation:
19     python run_phase3_eval.py --manifest ../manifest.json --score-only
20 """
21
22 import os
23 import sys
24 import json
25 import time
26 import argparse
27 import logging
28
29 from dotenv import load_dotenv
30 load_dotenv() # make GEMINI_API_KEY (etc.) available to the AI-backed segmenters
31
32 from backend.infrastructure.database import Database
33 from backend.config import load_config
34 from backend.utils.hashing import compute_video_id as get_file_md5 # canonical video_id
35 # Importing the package registers all segmenters (motion/gemini/yolo_hybrid/...).
36 from backend.pipeline.segmenters import segmenter_registry
37 from backend.eval.annotations import load_annotations
38 from backend.eval.harness import evaluate, format_report_text, report_to_dict
39 from backend.eval import batch
40
41 logger = logging.getLogger(__name__)
42
43
44 def build_parser() -> argparse.ArgumentParser:
45     p = argparse.ArgumentParser(
46         description="Batch process + evaluate a collector export manifest (Phase 3).",
47         formatter_class=argparse.RawDescriptionHelpFormatter,
48     )
49     p.add_argument("--manifest", required=True, help="Path to the collector's
manifest.json.")
50     p.add_argument("--videos-root", default=None,
51         help="Root for resolving manifest relative local_paths "
52         "(default: the collector repo root, two levels above the
manifest).")
53     p.add_argument("--segmenter", default="motion",
54         help="Segmenter to run (default: motion ? free/CPU). "
55         "yolo_hybrid/gemini need GEMINI_API_KEY and cost API calls.")
56     p.add_argument("--stage", choices=("candidates", "final", "both", "auto"),
default="auto",
57         help="Which prediction stage(s) to score (default: auto by
segmenter).")
58     p.add_argument("--detector", default=None, help="Restrict candidate scoring to this
detector.")
59     p.add_argument("--iou", type=float, default=0.5, help="IoU match threshold (default:
0.5).")
60     p.add_argument("--only", action="append", default=None,
61         help="Only this video_id (repeatable).")
62     p.add_argument("--limit", type=int, default=None, help="Process at most N videos.")
63     p.add_argument("--max-frames", type=int, default=None,
64         help="Cap segmenter frames (fast validation).")
65     p.add_argument("--score-only", action="store_true",
66         help="Skip segmentation; only score detections already in the DB.")
67     p.add_argument("--reprocess", action="store_true",
68         help="Re-run segmentation even if the video already has segments.")
69     p.add_argument("--db", default=None, help="Indexer DB path (default:
output/<db_name>).")
70     p.add_argument("--json", dest="json_out", default=None, help="Write a JSON report
here.")

```

```

71         p.add_argument("--show-failures", action="store_true", help="List missed/spurious
rallies per video.")
72         return p
73
74
75     def main(argv=None) -> int:
76         logging.basicConfig(level=logging.INFO, format="%(levelname)s %(message)s")
77         args = build_parser().parse_args(argv)
78
79         manifest_path = os.path.abspath(args.manifest)
80         if not os.path.exists(manifest_path):
81             print(f"ERROR: manifest not found: {manifest_path}", file=sys.stderr)
82             return 2
83         manifest_dir = os.path.dirname(manifest_path)
84         videos_root = args.videos_root or os.path.dirname(os.path.dirname(manifest_dir))
85
86         with open(manifest_path) as f:
87             manifest = json.load(f)
88             entries = manifest.get("eval_sets", [])
89
90         cfg = load_config() # typed AppConfig (same object the live pipeline feeds
segmenters)
91         db_path = args.db or os.path.join(cfg.output.output_dir, cfg.output.db_name)
92         os.makedirs(os.path.dirname(db_path) or ".", exist_ok=True)
93         db = Database(db_path)
94
95         stages = batch.default_stages_for(args.segmenter, args.stage)
96
97         # Resolve the segmenter class up front (unless purely scoring).
98         segmenter = None
99         if not args.score_only:
100             seg_cls = segmenter_registry.get(args.segmenter)
101             if not seg_cls:
102                 print(f"ERROR: segmenter '{args.segmenter}' not found. "
103                     f"Available: {list(segmenter_registry.list_available().keys())}",
file=sys.stderr)
104                 return 2
105             segmenter = seg_cls(db, cfg)
106
107         # Filter entries.
108         targets = []
109         for e in entries:
110             if args.only and e["video_id"] not in args.only:
111                 continue
112             targets.append(e)
113         if args.limit:
114             targets = targets[: args.limit]
115
116         per_video = []
117         skipped = []
118         print(f"Phase 3: {len(targets)} target(s) | segmenter={args.segmenter} |
stages={stages} | db={db_path}\n")
119
120         for e in targets:
121             vid = e["video_id"]
122             video_path = batch.resolve_video_path(e.get("local_path"), videos_root)
123             csv_path = os.path.join(manifest_dir, e["csv"])
124
125             if not video_path or not os.path.exists(video_path):
126                 skipped.append((vid, "no video file"))
127                 print(f"[skip] {vid}: no video file ({e.get('local_path')})")
128                 continue
129             if not os.path.exists(csv_path):
130                 skipped.append((vid, "no GT csv"))
131                 print(f"[skip] {vid}: GT csv missing ({csv_path})")

```

```

132         continue
133
134         video_id = get_file_md5(video_path)
135
136         # Segment unless score-only or already present (and not --reprocess).
137         if not args.score_only:
138             already = db.query_play_segments(video_id, {}) or
db.query_candidate_segments(video_id)
139             if already and not args.reprocess:
140                 print(f"[have] {vid}: detections already in DB (use --reprocess to
redo)")
141             else:
142                 print(f"[proc] {vid}: segmenting with {args.segmenter}...")
143                 db.add_video(video_id, video_path, "processed", [])
144                 t0 = time.time()
145                 segs = segmenter.process_video(video_id, video_path,
max_frames=args.max_frames)
146                 print(f"          -> {len(segs)} segments in {time.time() - t0:.0f}s")
147
148                 gt = load_annotations(csv_path)
149                 report = evaluate(db, video_id, gt, stages, iou_threshold=args.iou,
detector=args.detector)
150                 print(format_report_text(report, show_failures=args.show_failures))
151                 print()
152                 d = report_to_dict(report)
153                 d["collector_video_id"] = vid
154                 per_video.append(d)
155
156                 agg = batch.aggregate(per_video, stages)
157                 print(batch.format_aggregate(agg))
158                 if skipped:
159                     print(f"\nSkipped {len(skipped)}: " + ", ".join(f"{v} ({why})" for v, why in
skipped))
160
161                 if args.json_out:
162                     with open(args.json_out, "w") as f:
163                         json.dump({"segmenter": args.segmenter, "iou": args.iou, "stages": stages,
164                             "aggregate": agg, "per_video": per_video,
165                             "skipped": skipped}, f, indent=2)
166                     print(f"\nJSON report written to {args.json_out}")
167
168                 return 0
169
170
171 if __name__ == "__main__":
172     raise SystemExit(main())
173

```

29. user (2026-06-10T06:01:34.888Z)

```

1     """Load hand-annotated ground-truth rally timestamps from CSV.
2
3     Annotation file format (one file per video):
4
5         start,end
6         00:01:12,00:01:39
7         102.5,131.0
8
9     - A header row `start,end` is optional (auto-detected and skipped).
10    - Each timestamp is either decimal seconds (`102.5`) or colon-separated
11      `MM:SS` / `HH:MM:SS` (`00:01:12`). Both forms may be mixed in one file.
12    - Blank lines and `#` comment lines are ignored.
13    """
14
15    import csv

```

```

16 import logging
17 from typing import List, Tuple
18
19 logger = logging.getLogger(__name__)
20
21 Interval = Tuple[float, float]
22
23
24 def parse_timestamp(value: str) -> float:
25     """Parse a timestamp string to float seconds.
26
27     Accepts decimal seconds or colon-separated MM:SS / HH:MM:SS.
28     Raises ValueError on anything unparseable so bad annotations fail loudly.
29     """
30     value = value.strip()
31     if not value:
32         raise ValueError("empty timestamp")
33     if ":" in value:
34         parts = value.split(":")
35         if len(parts) > 3:
36             raise ValueError(f"too many ':' segments in timestamp '{value}'")
37         total = 0.0
38         for part in parts: # most-significant first (H, M, S)
39             total = total * 60 + float(part)
40         return total
41     return float(value)
42
43
44 def load_annotations(csv_path: str) -> List[Interval]:
45     """Load and validate ground-truth rally intervals from a CSV file.
46
47     Returns intervals sorted by start time. Raises ValueError on malformed rows
48     (bad timestamp, start >= end) so labeling mistakes surface immediately rather
49     than silently skewing scores.
50     """
51     intervals: List[Interval] = []
52     with open(csv_path, newline="") as f:
53         reader = csv.reader(f)
54         for lineno, row in enumerate(reader, start=1):
55             if not row:
56                 continue
57             first = row[0].strip()
58             if not first or first.startswith("#"):
59                 continue
60             # Skip an optional header row.
61             if first.lower() in ("start", "start_time"):
62                 continue
63             if len(row) < 2:
64                 raise ValueError(f"{csv_path}:{lineno}: expected 'start,end', got
65 {row!r}")
66             try:
67                 start = parse_timestamp(row[0])
68                 end = parse_timestamp(row[1])
69             except ValueError as e:
70                 raise ValueError(f"{csv_path}:{lineno}: {e}") from e
71             if start >= end:
72                 raise ValueError(f"{csv_path}:{lineno}: start ({start}) >= end ({end})")
73             intervals.append((start, end))
74
75     intervals.sort(key=lambda iv: iv[0])
76     logger.info("Loaded %d ground-truth rallies from %s", len(intervals), csv_path)
77     return intervals
78
79 def write_scaffold(csv_path: str) -> None:

```

```

80         """Write an empty annotation file with just the header, for hand-filling."""
81         with open(csv_path, "w", newline="") as f:
82             writer = csv.writer(f)
83             writer.writerow(["start", "end"])
84

```

30. user (2026-06-10T06:01:47.346Z)

```

28-Interval = Tuple[float, float]
29-Config = Tuple[float, float, float] # (velocity_thresh, merge_gap, min_window_duration)
30-
31-BASELINE: Config = (0.02, 2.0, 2.0) # the WasbConfig / trajectory_to_action_windows
defaults
32-
33-
34:def load_golden_gts(csv_path: str) -> List[Interval]:
35:    """Rally [start,end] seconds from the golden CSV (start_time/end_time cols)."""
36:    gts: List[Interval] = []
37:    with open(csv_path, newline="", encoding="utf-8") as f:
38:        reader = csv.DictReader(f)
39:        reader.fieldnames = [(n or "").strip().lower() for n in (reader.fieldnames or [])]
40:        for row in reader:
41:            row = {(k or "").strip().lower(): v for k, v in row.items()}
42:            try:
43:                s = float((row.get("start_time") or "").strip())
44:                e = float((row.get("end_time") or "").strip())
45:            except (TypeError, ValueError):
46:                continue
47:            if e > s:
48:                gts.append((s, e))
49:    return sorted(gts)
--
100-    return {"baseline": base.as_dict(), "best_cfg": best_cfg, "fit": fit, "cv": cv}
101-
102-
103-def main():
104-    ap = argparse.ArgumentParser()
105-    ap.add_argument("--trajectory", required=True, help="WASB trajectory CSV
(Frame,Visibility,X,Y)")
106-    ap.add_argument("--labels", required=True, help="golden labels CSV (start_time/end_time
cols)")
107-    ap.add_argument("--fps", type=float, default=59.94)
108-    ap.add_argument("--frame-width", type=float, default=1920.0)
109-    ap.add_argument("--iou", type=float, default=0.5)
110-    args = ap.parse_args()
111-    run(args.trajectory, args.labels, args.fps, args.frame_width, args.iou)
112-

```

31. user (2026-06-10T06:01:48.809Z)

```

6-simply un-watched (not "no rally there"). Scoring naively against such labels is wrong:
7-the detector's *correct* rallies in the un-labeled tail would be counted as false
8-positives and crush precision.
9-
10-This harness fixes that by restricting the evaluation to the LABELED WINDOW
11:[window_start, window_end] (window_end defaults to the last labeled rally's end) and
12-dropping predicted segments that fall outside it. Within that window the human labels
13-ARE complete, so precision/recall/F1 are honest.
14-
15-It reuses everything from the existing eval stack (no new windowing\metrics\gate logic):
16- - calibrate_wasb.make_predict_fn / load_golden_gts / BASELINE
--
42-# Operating points to compare. The tuned config wins on its fit-video but over-segments
43-# and does not transfer; baseline (+ net-crossing gate) is the stronger general point.
44-TUNED: Config = (0.05, 1.0, 1.0)

```

```

45-
46-
47: def restrict_to_window(intervals: List[Interval], w0: float, w1: float) -> List[Interval]:
48-     """Keep intervals that overlap the labeled window [w0, w1].
49-
50-     A predicted segment is evaluated iff it overlaps the labeled region; segments wholly
51-     in the un-labeled tail (start >= w1) or before the window (end <= w0) are dropped so
52-     they are neither rewarded nor penalized. GTs are assumed already inside the window.
53-     --
64-     sc = score(preds, gts, iou)
65-     return {"cfg": cfg, "gate": gate_crossings, "n_pred": len(preds), "score": sc, "preds":
preds}
66-
67-
68: def run(trajjectory_csv: str, labels_csv: str, fps: float, frame_width: float,
69-         window_start: float = 0.0, window_end: Optional[float] = None,
70-         gate_crossings: int = 2, iou: float = 0.5,
71-         comparison_out: Optional[str] = None) -> dict:
72-     gts_all = load_golden_gts(labels_csv)
73-     if not gts_all:
74-         raise SystemExit(f"no usable human rallies in {labels_csv}")
75-     if window_end is None:
76-         window_end = max(e for _, e in gts_all)
77-     window = (window_start, window_end)
78-     gts = restrict_to_window(gts_all, *window)
79-
80-     predict_fn = make_predict_fn(trajjectory_csv, fps, frame_width)
81-     points = TrackNetRunner.parse_trajectory_csv(trajjectory_csv)
82-     traj_end = (max(p.frame for p in points) / fps) if points else 0.0
83-
84-     print("=== Pipeline vs PARTIAL human labels ===")
85-     print(f"labels: {labels_csv}")
86-     print(f" {len(gts_all)} human rallies; labeled window [{window_start:.1f}s,
{window_end:.1f}s] "
87-           f"-> {len(gts)} rallies in-window")
88-     print(f"trajectory covers ~0-{traj_end:.1f}s "
89-           f"({'covers the window' if traj_end >= window_end - 1 else 'WARNING: shorter than
window'})")
90-     print(f"IoU>={iou}; net-crossing gate K={gate_crossings}\n")
91-
92-     # Compare the standard operating points: baseline / baseline+gate / tuned / tuned+gate.
93-     points_for_gate = points
94-     candidates = [
95-         --
109-         results.append(r)
110-         print(f"{name:<18} {r['n_pred']:>4} {sc.true_positives:>3} {sc.false_negatives:>3} "
111-               f"{sc.false_positives:>3} {sc.precision:>6.3f} {sc.recall:>6.3f} {sc.f1:>6.3f} "
112-               f"{sc.mean_iou:>6.3f} {sc.mean_start_error:>6.2f} {sc.mean_end_error:>6.2f}")
113-
114-     best = max(results, key=lambda r: r["score"].f1)
115-     print(f"\nbest operating point by F1: {best['name']} (F1 {best['score'].f1:.3f})")
116-
117-     # Per-rally detail for the best point: which human rallies matched/missed + extra fires.
118-     comp = build_comparison(best["preds"], gts, iou)
119-     comp_out = comparison_out or osp.join("output", "adarsh_avi_vs_human.csv")
120-     --
138-     for mwd in (0.5, 1.0, 1.5, 2.0)
139- ]
140-
141-
142: def sweep(trajjectory_csv: str, labels_csv: str, fps: float, frame_width: float,
143-          window_start: float = 0.0, window_end: Optional[float] = None,
144-          gate_options: Tuple[int, ...] = (0, 2), iou: float = 0.5,
145-          grid: Optional[List[Config]] = None, top: int = 12) -> dict:

```

```

146-     """Grid-search (velocity, merge_gap, min_dur) x gate vs human labels, in-window.
147-
148- WARNING: this is fit-on-self on ONE video -> optimistic. It shows the achievable
--
151-     """
152-     grid = grid or DEFAULT_GRID
153-     gts_all = load_golden_gts(labels_csv)
154-     if not gts_all:
155-         raise SystemExit(f"no usable human rallies in {labels_csv}")
156-     if window_end is None:
157-         window_end = max(e for _, e in gts_all)
158-     window = (window_start, window_end)
159-     gts = restrict_to_window(gts_all, *window)
160-     predict_fn = make_predict_fn(trajjectory_csv, fps, frame_width)
161-     points = TrackNetRunner.parse_trajectory_csv(trajjectory_csv)
162-
163-     rows: List[dict] = []
--
169-         "precision": sc.precision, "recall": sc.recall, "mean_iou":
sc.mean_iou,
170-         "tp": sc.true_positives, "fp": sc.false_positives, "fn":
sc.false_negatives})
171-     rows.sort(key=lambda x: (x["f1"], x["mean_iou"]), reverse=True)
172-
173-     print(f"=== Config sweep vs human labels ({len(gts)} rallies in "
174-           f"[{window_start:.0f},{window_end:.0f}]s; {len(grid)}x{len(gate_options)} configs)
===")
175-     print("NOTE: fit-on-self on ONE video = optimistic; real test is leave-one-video-
out.\n")
176-     hdr = f"{'#':>3} {'veloc':>6} {'merge':>6} {'minDur':>6} {'gate':>4} " \
177-           f"{'F1':>6} {'prec':>6} {'recall':>6} {'mIoU':>6} {'TP':>3} {'FP':>3} {'FN':>3}"
178-     print(hdr); print("-" * len(hdr))
179-     for i, r in enumerate(rows[:top], 1):
--
203-         help="grid-search configs vs labels (fit-on-self; finds the ceiling +
best levers)")
204-     ap.add_argument("--top", type=int, default=12, help="how many sweep rows to print")
205-     args = ap.parse_args()
206-     if args.sweep:
207-         sweep(args.trajectory, args.labels, args.fps, args.frame_width,
208-               window_start=args.window_start, window_end=args.window_end,
209-               iou=args.iou, top=args.top)
210-     else:
211-         run(args.trajectory, args.labels, args.fps, args.frame_width,
212-             window_start=args.window_start, window_end=args.window_end,
213-             gate_crossings=args.gate, iou=args.iou, comparison_out=args.comparison_out)
214-
215-
216-if __name__ == "__main__":
217-     main()

```

32. user (2026-06-10T06:02:45.870Z)

```

1     import argparse
2     import sys
3     from src.core.config import load_config, get_database_path
4     from src.core.db import SportsDatabase
5     from src.parsers.badminton.coachai import CoachAIBadmintonParser
6
7     def main():
8         parser = argparse.ArgumentParser(description="Ingest external datasets into the
curation staging area")
9         parser.add_argument("--sport", required=True, choices=["badminton"], help="Sport
type")
10        parser.add_argument("--parser", required=True, choices=["coachai"], help="Parser

```

```

type to use")
11     parser.add_argument("--file-path", required=True, help="Path to raw dataset file")
12     parser.add_argument("--license", default="MIT", help="Dataset license type")
13     parser.add_argument("--video-url", default="https://youtube.com/watch?v=example",
help="URL to the video")
14     parser.add_argument("--match-name", help="Name of the match")
15     parser.add_argument("--video-id", help="Video unique identifier")
16
17     args = parser.parse_args()
18
19     # Load db path from the central config (falls back to the default).
20     # NOTE: this importer is intentionally badminton/CoachAI-only for now; other
21     # sports are populated via `curate import-local` or vendor golden sets.
22     db = SportsDatabase(get_database_path(load_config()))
23
24     if args.sport == "badminton" and args.parser == "coachai":
25         match_parser = CoachAIBadmintonParser()
26         try:
27             video_meta, rallies = match_parser.parse(
28                 args.file_path,
29                 video_id=args.video_id,
30                 video_url=args.video_url,
31                 license=args.license,
32                 match_name=args.match_name
33             )
34
35             # Insert into database (by default parsed entries are marked as staging)
36             db.insert_video(video_meta)
37             count = db.insert_badminton_rallies(rallies)
38
39             print(f"Successfully ingested match '{video_meta.match_name}' into
staging:")
40             print(f" - Video ID: {video_meta.video_id}")
41             print(f" - Total Rallies: {count}")
42             print(f" - License: {video_meta.license_type.value}")
43             print(f" - Commercial safe: {video_meta.is_commercial}")
44         except Exception as e:
45             print(f"Error parsing file: {e}")
46             sys.exit(1)
47
48     if __name__ == "__main__":
49         main()
50

```

33. user (2026-06-10T06:02:46.419Z)

```

1     """Per-video observability reports for the collector.
2
3     `obsreport` is a SOFT dependency: absent -> these emitters no-op and the CLI runs
4     unchanged. Every emitter is safe to call from normal control flow and NEVER raises
5     (emitting a report must not break ingestion). Returns the written paths, or None.
6     """
7     from __future__ import annotations
8
9     import os
10    from typing import Any, Dict, List, Optional
11
12    try: # soft dependency
13        import obsreport # noqa: F401
14        from . import adapter
15
16        OBSREPORT_AVAILABLE = True
17    except Exception: # pragma: no cover - only when obsreport is absent
18        OBSREPORT_AVAILABLE = False
19

```



```

20
21 def emit_delivery_report(
22     *,
23     out_dir: str,
24     source: str,
25     records: List[Any],
26     issues: List[Any],
27     fps: Optional[float],
28     video_path: Optional[str] = None,
29     video_duration_s: Optional[float] = None,
30     source_format: str = "cvat",
31     step: Optional[int] = None,
32 ) -> Optional[Dict[str, str]]:
33     """Write the delivery observability report into out_dir (alongside the existing
34     <id>_report.* artifacts). Used by validate-delivery / convert-cvat. Never raises."""
35     if not OBSREPORT_AVAILABLE:
36         return None
37     try:
38         rec = adapter.build_delivery_recorder(
39             records=records, issues=issues, fps=fps, source=source,
40             video_path=video_path,
41             video_duration_s=video_duration_s, source_format=source_format, step=step,
42         )
43         res = rec.write(output_dir=out_dir, stem=source)
44         return res.paths or None
45     except Exception: # pragma: no cover - reporting must never break the command
46         return None
47
48 def emit_import_report(
49     *,
50     video_meta_obj: Any,
51     annotations: List[Any],
52     is_valid: bool,
53     errors: List[str],
54     count: int,
55     sport: str,
56     source_format: str,
57     fallback_dir: Optional[str] = None,
58 ) -> Optional[Dict[str, str]]:
59     """Write the import observability report next to the imported video. If the video
60     file doesn't exist on disk (registration-only imports), fall back to
61     ``fallback_dir`` (e.g. the annotations directory) rather than the cwd. Never
62     raises."""
63     if not OBSREPORT_AVAILABLE:
64         return None
65     try:
66         rec = adapter.build_import_recorder(
67             video_meta_obj=video_meta_obj, annotations=annotations, is_valid=is_valid,
68             errors=errors, count=count, sport=sport, source_format=source_format,
69         )
70         video_path = getattr(video_meta_obj, "local_path", None)
71         if video_path and os.path.exists(video_path):
72             res = rec.write() # next to the video
73         else:
74             res = rec.write(output_dir=fallback_dir or ".",
75                             stem=getattr(video_meta_obj, "video_id", None))
76         return res.paths or None
77     except Exception: # pragma: no cover
78         return None

```

34. user (2026-06-10T06:02:47.136Z)

```
1      #!/usr/bin/env python3
```

```

2     import os
3     import sys
4     import shutil
5     import argparse
6     import math
7
8     from src.core.config import load_config
9
10    def get_file_or_dir_size(path):
11        """Calculates the size of a file or directory in bytes."""
12        if not os.path.exists(path):
13            return 0
14        if os.path.isfile(path):
15            return os.path.getsize(path)
16
17        total_size = 0
18        for dirpath, _, filenames in os.walk(path):
19            for f in filenames:
20                fp = os.path.join(dirpath, f)
21                if os.path.exists(fp) and not os.path.islink(fp):
22                    total_size += os.path.getsize(fp)
23        return total_size
24
25    def format_size(size_bytes):
26        """Formats bytes to a human-readable string."""
27        if size_bytes == 0:
28            return "0 B"
29        size_name = ("B", "KB", "MB", "GB")
30        i = int(math.floor(math.log(size_bytes, 1024)))
31        p = math.pow(1024, i)
32        s = round(size_bytes / p, 2)
33        return f"{s} {size_name[i]}"
34
35    def get_all_items_in_dir(path):
36        """Returns a list of absolute paths of all files/directories inside a directory."""
37        items = []
38        if not os.path.exists(path) or not os.path.isdir(path):
39            return items
40        for entry in os.listdir(path):
41            items.append(os.path.join(path, entry))
42        return items
43
44    def main():
45        parser = argparse.ArgumentParser(
46            description="Purge stored state including SQLite database, downloaded videos,
and temporary files."
47        )
48        parser.add_argument(
49            "--config",
50            default="config/config.yaml",
51            help="Path to the config.yaml file (default: config/config.yaml)"
52        )
53        parser.add_argument(
54            "--yes", "-y",
55            action="store_true",
56            help="Bypass confirmation prompt and proceed directly with deletion"
57        )
58        parser.add_argument(
59            "--dry-run",
60            action="store_true",
61            help="Preview what would be deleted without making any changes"
62        )
63
64        args = parser.parse_args()
65

```

```

66     # Load configuration
67     config_path = args.config
68     db_path = "data/sports_metadata.db"
69     video_dir = "data/videos"
70     temp_dir = "data/temp"
71
72     config = load_config(config_path)
73     if config:
74         print(f"Loaded configuration from: {config_path}")
75         db_path = config.get("database_path", db_path)
76         storage = config.get("video_storage", {})
77         video_dir = storage.get("local_dir", video_dir)
78     else:
79         print(f"No usable config at '{config_path}'. Using default paths.")
80
81     # Normalize paths
82     db_path = os.path.abspath(db_path)
83     video_dir = os.path.abspath(video_dir)
84     temp_dir = os.path.abspath(temp_dir)
85
86     print("\n--- Cleanup Target Summary ---")
87
88     # 1. Database
89     db_exists = os.path.exists(db_path)
90     db_size = get_file_or_dir_size(db_path) if db_exists else 0
91     print(f"Database File : {db_path}")
92     print(f" - Status: {'Exists' if db_exists else 'Not Found'}")
93     if db_exists:
94         print(f" - Size : {format_size(db_size)}")
95
96     # 2. Videos directory
97     video_files = get_all_items_in_dir(video_dir)
98     video_dir_size = get_file_or_dir_size(video_dir)
99     print(f"Videos Directory: {video_dir}")
100    print(f" - Status: {'Exists' if os.path.exists(video_dir) else 'Not Found'}")
101    print(f" - Files found: {len(video_files)}")
102    if video_files:
103        print(f" - Total Size : {format_size(video_dir_size)}")
104
105    # 3. Temp directory
106    temp_files = get_all_items_in_dir(temp_dir)
107    temp_dir_size = get_file_or_dir_size(temp_dir)
108    print(f"Temp Directory : {temp_dir}")
109    print(f" - Status: {'Exists' if os.path.exists(temp_dir) else 'Not Found'}")
110    print(f" - Files found: {len(temp_files)}")
111    if temp_files:
112        print(f" - Total Size : {format_size(temp_dir_size)}")
113
114    total_reclaimable = db_size + video_dir_size + temp_dir_size
115    print("-----")
116    print(f"Total space reclaimable: {format_size(total_reclaimable)}\n")
117
118    # If there is nothing to clean up
119    nothing_to_do = not db_exists and not video_files and not temp_files
120    if nothing_to_do:
121        print("All targets are already empty. Nothing to clean up.")
122        return 0
123
124    if args.dry_run:
125        print("Dry-run mode active. No changes will be written to disk.")
126        print("Would perform the following operations:")
127        if db_exists:
128            print(f" [DELETE] {db_path}")
129        if video_files:
130            print(f" [PURGE] {video_dir} ({len(video_files)} items)")

```

```

131         if temp_files:
132             print(f" [PURGE] {temp_dir} ({len(temp_files)} items)")
133             print("\nDry-run complete.")
134             return 0
135
136     # Confirmation
137     if not args.yes:
138         confirm = input("Are you sure you want to delete these files and clear the
stored state? [y/N]: ").strip().lower()
139         if confirm not in ("y", "yes"):
140             print("Aborted. No files were deleted.")
141             return 0
142
143     print("\nExecuting cleanup...")
144
145     # Delete Database
146     if db_exists:
147         try:
148             os.remove(db_path)
149             print(f"Deleted database file: {db_path}")
150         except Exception as e:
151             print(f"Error deleting database file: {e}")
152
153     # Purge Videos Directory
154     if os.path.exists(video_dir):
155         print(f"Purging video storage directory: {video_dir}")
156         for item in video_files:
157             try:
158                 if os.path.isdir(item):
159                     shutil.rmtree(item)
160                 else:
161                     os.remove(item)
162             except Exception as e:
163                 print(f" - Error removing {os.path.basename(item)}: {e}")
164     else:
165         # Recreate if it doesn't exist
166         try:
167             os.makedirs(video_dir, exist_ok=True)
168             print(f"Created video storage directory: {video_dir}")
169         except Exception as e:
170             print(f"Error creating video storage directory: {e}")
171
172     # Purge Temp Directory
173     if os.path.exists(temp_dir):
174         print(f"Purging temp directory: {temp_dir}")
175         for item in temp_files:
176             try:
177                 if os.path.isdir(item):
178                     shutil.rmtree(item)
179                 else:
180                     os.remove(item)
181             except Exception as e:
182                 print(f" - Error removing {os.path.basename(item)}: {e}")
183     else:
184         # Recreate if it doesn't exist
185         try:
186             os.makedirs(temp_dir, exist_ok=True)
187             print(f"Created temp directory: {temp_dir}")
188         except Exception as e:
189             print(f"Error creating temp directory: {e}")
190
191     print("\nCleanup complete! All targets have been purged.")
192     return 0
193
194 if __name__ == "__main__":

```

```
195         sys.exit(main())
196
```

35. user (2026-06-10T06:02:47.257Z)

```
1     [build-system]
2     requires = ["setuptools>=61.0.0"]
3     build-backend = "setuptools.build_meta"
4
5     [project]
6     name = "sports-data-collector"
7     version = "0.1.0"
8     description = "A collector and curator for sports video datasets with sport-specific
schemas and license checks."
9     readme = "README.md"
10    requires-python = ">=3.9"
11    dependencies = [
12        "pydantic>=2.0.0",
13        "pyyaml>=6.0",
14        "tabulate>=0.9.0"
15    ]
16
17    [project.optional-dependencies]
18    # Video acquisition (`fetch` / crawler). Also requires the external binaries
19    # ffprobe (FFmpeg) and, for YouTube high-res, a JS runtime such as deno ?
20    # neither is pip-installable; see README "External requirements".
21    fetch = [
22        "yt-dlp>=2024.1.0"
23    ]
24    dev = [
25        "pytest>=7.0.0",
26        "pytest-cov>=4.0.0"
27    ]
28    # Per-video observability reports (SOFT dependency: ingestion degrades to a no-op
29    # if absent). Shared, permissively-licensed report toolkit. Pin a tag once the
30    # private repo is published, e.g.:
31    #   obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3
32    # For local development install it editable from a sibling checkout:
33    #   pip install -e ../sports-obsreport --no-build-isolation
34    report = [
35        "obsreport>=0.1.3",
36    ]
37
```

36. user (2026-06-10T06:02:58.026Z)

```
1     """End-to-end coverage for the SportsDataCLI commands.
2
3     Each test runs against a throwaway config + SQLite DB in tmp_path, so the real
4     data/ directory is never touched. Network (yt-dlp) and the browser are mocked.
5     """
6     import csv as _csv
7     import json
8     from types import SimpleNamespace
9     from unittest import mock
10
11    import pytest
12    import yaml
13
14    from src.cli.curate import SportsDataCLI
15    from src.core.schemas.base import VideoMetadata, SportType, LicenseType, CurationStatus
16    from src.core.schemas.badminton import BadmintonRally
17
18
19    @pytest.fixture
```

```

20     def cli(tmp_path):
21         cfg = {
22             "database_path": str(tmp_path / "test.db"),
23             "video_storage": {
24                 "local_dir": str(tmp_path / "videos"),
25                 "max_download_size_mb": 100,
26                 "auto_download": False,
27                 "max_height": None,
28             },
29             "commercial_licenses": ["MIT", "Apache-2.0", "BSD", "CC0", "CC-BY", "Public-
Domain", "Proprietary"],
30             "supported_sports": ["badminton", "tennis", "tabletennis", "pickleball",
"cricket"],
31         }
32         cfg_path = tmp_path / "config.yaml"
33         cfg_path.write_text(yaml.safe_dump(cfg), encoding="utf-8")
34         return SportsDataCLI(config_path=str(cfg_path))
35
36
37     def _import_args(**kw):
38         base = dict(
39             sport=None, video_path="nonexistent.mp4", annotations_path=None,
40             match_name=None, video_id=None, license="Proprietary", fps=30.0,
resolution="1080p",
41         )
42         base.update(kw)
43         return SimpleNamespace(**base)
44
45
46     def _write_csv(path, header, rows):
47         with open(path, "w", newline="", encoding="utf-8") as f:
48             w = _csv.writer(f)
49             w.writerow(header)
50             for r in rows:
51                 w.writerow(r)
52
53
54     # ----- #
55     # import-local: every sport, both JSON and CSV (covers dispatch + parse) #
56     # ----- #
57
58     def test_import_local_badminton_json(cli, tmp_path):
59         p = tmp_path / "bad.json"
60         p.write_text(json.dumps([
61             {"rally_number": 1, "start_time": 10.0, "end_time": 20.0, "shots_count": 5,
"ending_reason": "winner"},
62             {"rally_number": 2, "start_time": 25.0, "end_time": 30.0, "shots_count": 1,
"ending_reason": "service_fault"},
63         ]), encoding="utf-8")
64         cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
video_id="bad1", match_name="Bad"))
65         assert len(cli.db.get_badminton_rallies("bad1")) == 2
66
67
68     def test_import_local_pickleball_csv(cli, tmp_path):
69         p = tmp_path / "pb.csv"
70         _write_csv(p, ["rally_number", "start_time", "end_time", "shots_count",
"ending_reason"],
71             [[1, 10.0, 20.0, 6, "winner"], [2, 25.0, 27.0, 1, "service_fault"]])
72         cli.import_local(_import_args(sport="pickleball", annotations_path=str(p),
video_id="pb1", match_name="PB"))
73         assert len(cli.db.get_pickleball_rallies("pb1")) == 2
74
75
76     def test_import_local_tabletennis_json(cli, tmp_path):

```

```

77         p = tmp_path / "tt.json"
78         p.write_text(json.dumps([
79             {"rally_number": 1, "start_time": 5.0, "end_time": 9.0, "shots_count": 7,
"ending_reason": "winner"},
80         ]), encoding="utf-8")
81         cli.import_local(_import_args(sport="tabletennis", annotations_path=str(p),
video_id="tt1", match_name="TT"))
82         assert len(cli.db.get_tabletennis_rallies("tt1")) == 1
83
84
85     def test_import_local_tennis_csv(cli, tmp_path):
86         p = tmp_path / "tn.csv"
87         _write_csv(p, [{"point_number", "start_time", "end_time", "server", "ending_type"},
88             [[1, 10.0, 20.0, "player1", "winner"]]])
89         cli.import_local(_import_args(sport="tennis", annotations_path=str(p),
video_id="tn1", match_name="TN"))
90         assert len(cli.db.get_tennis_points("tn1")) == 1
91
92
93     def test_import_local_cricket_json(cli, tmp_path):
94         p = tmp_path / "ck.json"
95         p.write_text(json.dumps([
96             {"over": 0, "ball": 1, "start_time": 5.0, "end_time": 15.0, "runs": 4, "wicket":
False},
97         ]), encoding="utf-8")
98         cli.import_local(_import_args(sport="cricket", annotations_path=str(p),
video_id="ck1", match_name="CK"))
99         assert len(cli.db.get_cricket_deliveries("ck1")) == 1
100
101
102     # ----- #
103     # import-local: error / edge branches #
104     # ----- #
105
106     def test_import_local_unsupported_sport_exits(cli, tmp_path):
107         p = tmp_path / "x.json"
108         p.write_text("[]", encoding="utf-8")
109         with pytest.raises(SystemExit):
110             cli.import_local(_import_args(sport="football", annotations_path=str(p)))
111
112
113     def test_import_local_missing_annotations_exits(cli):
114         with pytest.raises(SystemExit):
115             cli.import_local(_import_args(sport="badminton",
annotations_path="does_not_exist.json"))
116
117
118     def test_import_local_overlap_fails_validation(cli, tmp_path):
119         p = tmp_path / "ov.csv"
120         _write_csv(p, [{"rally_number", "start_time", "end_time"},
121             [[1, 10.0, 20.0], [2, 15.0, 25.0]]]) # overlap
122         with pytest.raises(SystemExit):
123             cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
video_id="ov1"))
124
125
126     def test_import_local_unknown_license_defaults(cli, tmp_path):
127         p = tmp_path / "lic.json"
128         p.write_text(json.dumps([{"rally_number": 1, "start_time": 1.0, "end_time": 2.0}],
encoding="utf-8"))
129         cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
video_id="lic1", license="WeirdLicense"))
130         v = cli.db.get_video("lic1")
131         assert v.license_type == LicenseType.UNKNOWN # unknown string coerces to Unknown
132         assert v.is_commercial is False # Unknown is non-commercial

```

```

133
134
135     def test_import_local_json_missing_required_field_exits(cli, tmp_path):
136         # No start_time -> hardened JSON path raises a clear ValueError -> SystemExit
137         p = tmp_path / "missing.json"
138         p.write_text(json.dumps([{"rally_number": 1, "end_time": 5.0}], encoding="utf-8"))
139         with pytest.raises(SystemExit):
140             cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
141 video_id="miss1"))
142
143     def test_import_local_bad_extension_exits(cli, tmp_path):
144         p = tmp_path / "bad.txt"
145         p.write_text("nope", encoding="utf-8")
146         with pytest.raises(SystemExit):
147             cli.import_local(_import_args(sport="badminton", annotations_path=str(p)))
148
149
150     # ----- #
151     # status + export                                     #
152     # ----- #
153
154     def test_status_empty_and_populated(cli, tmp_path):
155         cli.status() # empty branch
156         p = tmp_path / "b.json"
157         p.write_text(json.dumps([{"rally_number": 1, "start_time": 1.0, "end_time": 2.0}],
158 encoding="utf-8"))
159         cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
160 video_id="s1"))
161         cli.status() # populated branch
162
163
164     def test_export_eval_set(cli, tmp_path):
165         p = tmp_path / "b.json"
166         p.write_text(json.dumps([
167             {"rally_number": 1, "start_time": 10.0, "end_time": 20.0},
168             {"rally_number": 2, "start_time": 25.0, "end_time": 35.0},
169         ]), encoding="utf-8")
170         cli.import_local(_import_args(sport="badminton", annotations_path=str(p),
171 video_id="exp1", license="MIT"))
172
173         out_dir = tmp_path / "out"
174         args = SimpleNamespace(out_dir=str(out_dir), sport=None, include_staging=False,
175 include_noncommercial=False)
176         cli.export_eval_set(args)
177
178         manifest = out_dir / "manifest.json"
179         assert manifest.exists()
180         doc = json.loads(manifest.read_text(encoding="utf-8"))
181         assert doc["total_videos"] == 1
182         assert doc["total_segments"] == 2
183         assert (out_dir / "badminton" / "exp1.csv").exists()
184
185
186     # ----- #
187     # interactive curate loop (mocked input + browser)         #
188     # ----- #
189
190     def _stage_badminton(cli, vid="stg1"):
191         cli.db.insert_video(VideoMetadata(
192             video_id=vid, sport=SportType.BADMINTON, license_type=LicenseType.MIT,
193             is_commercial=True, match_name="Staged Match", status=CurationStatus.STAGING,
194             video_url="https://www.youtube.com/watch?v=abc123",
195         ))
196         cli.db.insert_badminton_rallies([

```



```

193         BadmintonRally(video_id=vid, rally_number=1, start_time=10.0, end_time=20.0,
shots_count=5),
194     ])
195
196
197     def test_curate_quit(cli):
198         _stage_badminton(cli)
199         with mock.patch("builtins.input", side_effect=["q"]):
200             cli.curate()
201
202
203     def test_curate_empty(cli):
204         cli.curate() # no staging entries branch
205
206
207     def test_curate_review_keep(cli):
208         _stage_badminton(cli)
209         with mock.patch("builtins.input", side_effect=["1", "k"]):
210             cli.curate()
211         assert cli.db.get_video("stg1").status == CurationStatus.STAGING
212
213
214     def test_curate_approve(cli):
215         _stage_badminton(cli)
216         with mock.patch("builtins.input", side_effect=["1", "a"]):
217             cli.curate()
218         assert cli.db.get_video("stg1").status == CurationStatus.CURATED
219
220
221     def test_curate_reject(cli):
222         _stage_badminton(cli)
223         with mock.patch("builtins.input", side_effect=["1", "r"]):
224             cli.curate()
225         assert cli.db.get_video("stg1") is None
226
227
228     def test_curate_play_segment(cli):
229         _stage_badminton(cli)
230         with mock.patch("builtins.input", side_effect=["1", "p", "1", "k"]), \
231             mock.patch("webbrowser.open") as wb:
232             cli.curate()
233         wb.assert_called_once()
234
235
236     # ----- #
237     # fetch (mocked downloader) #
238     # ----- #
239
240     def _fetch_args(**kw):
241         base = dict(
242             sport=None, video_id=None, status="all", min_height=720, max_height=None,
243             force=False, include_noncommercial=False, cookies_from_browser=None,
244             cookies=None, player_client="default", sleep=0.0, retries=0, retry_backoff=0.0,
245         )
246         base.update(kw)
247         return SimpleNamespace(**base)
248
249
250     def test_fetch_downloads_and_updates_db(cli, tmp_path):
251         _stage_badminton(cli, "fvid")
252         fake = SimpleNamespace(
253             success=True, height=1080, path=str(tmp_path / "videos" / "fvid.mp4"),
254             resolution_label="1080p", fps=30.0, size_mb=50.0, error=None,
255         )
256         with mock.patch("src.cli.curate.download_best_video", return_value=fake):

```

```

257         cli.fetch(_fetch_args())
258         assert cli.db.get_video("fvid").local_path == fake.path
259
260
261     def test_fetch_handles_failure(cli):
262         _stage_badminton(cli, "fvid2")
263         fake = SimpleNamespace(
264             success=False, height=None, path=None, resolution_label=None,
265             fps=None, size_mb=0.0, error="boom",
266         )
267         with mock.patch("src.cli.curate.download_best_video", return_value=fake):
268             cli.fetch(_fetch_args())
269             # Failed fetch leaves the row unchanged (no local_path written).
270             assert cli.db.get_video("fvid2").local_path is None
271
272
273     # ----- #
274     # main() argparse dispatch #
275     # ----- #
276
277     @pytest.mark.parametrize("argv,method", [
278         (["prog", "status"], "status"),
279         (["prog", "curate"], "curate"),
280         (["prog", "export"], "export_eval_set"),
281         (["prog", "fetch"], "fetch"),
282     ])
283     def test_main_dispatch(argv, method):
284         from src.cli import curate as curate_mod
285         with mock.patch.object(curate_mod, "SportsDataCLI") as M, mock.patch("sys.argv",
argv):
286             curate_mod.main()
287             getattr(M.return_value, method).assert_called_once()
288
289
290     def test_main_import_local_dispatch():
291         from src.cli import curate as curate_mod
292         argv = ["prog", "import-local", "--sport", "badminton", "--video-path", "v.mp4", "--
annotations-path", "a.json"]
293         with mock.patch.object(curate_mod, "SportsDataCLI") as M, mock.patch("sys.argv",
argv):
294             curate_mod.main()
295             M.return_value.import_local.assert_called_once()
296
297
298     def test_main_no_command_prints_help():
299         from src.cli import curate as curate_mod
300         with mock.patch.object(curate_mod, "SportsDataCLI"), mock.patch("sys.argv",
["prog"]):
301             curate_mod.main() # falls through to print_help()
302

```

37. user (2026-06-10T06:02:58.296Z)

```

1     """Tests for the CVAT rally adapter (src/parsers/cvat/rally_cvat.py).
2
3     The adapter's whole reason for existing is to (a) recompute rally time from the
4     box FRAME RANGE at the true fps ? ignoring the vendor's wrong start/end_time
5     attributes ? and (b) report content defects rather than silently fixing them.
6     These tests pin both behaviours, plus the end-to-end hand-off to import-local.
7     """
8     import csv as _csv
9
10    from src.parsers.cvat.rally_cvat import (
11        RallyRecord, parse_cvat_boxes, build_records, detect_issues, convert,
12        write_canonical_csv, load_canonical_csv, build_summary, render_validation_md,

```

```

13         write_review_csv,
14     )
15     from src.cli.curate import SportsDataCLI
16
17
18     # --- helpers -----
19
20     def _img_box(frame, attrs):
21         rows = "".join(f'<attribute name="{k}">{v}</attribute>' for k, v in attrs.items())
22         return (f'<image id="{frame}" name="frame_{frame:06d}" width="1920" height="1080">'
23             f'<box label="rally" xtl="1" ytl="1" xbr="2" ybr="2">{rows}</box></image>')
24
25
26     def _write_cvat(path, images, step=30, stop_frame=1000):
27         body = "".join(images)
28         xml = (
29             '<?xml version="1.0" encoding="utf-8"?><annotations><version>1.1</version>'
30             f'<meta><job><stop_frame>{stop_frame}</stop_frame>'
31             f'<frame_filter>step={step}</frame_filter></job></meta>'
32             f'{body}</annotations>'
33         )
34         path.write_text(xml, encoding="utf-8")
35         return str(path)
36
37
38     def _rec(rn, start, end, shots=5, reason="winner", fmin=None, fmax=None, n=3,
39 outcome=None):
40         return RallyRecord(
41             rally_number=rn, start_time=start, end_time=end, shots_count=shots,
42             ending_reason=reason, last_shot_outcome=outcome,
43             score_before=None, score_after=None, notes=None,
44             frame_min=fmin if fmin is not None else int(start * 60),
45             frame_max=fmax if fmax is not None else int(end * 60), n_boxes=n)
46
47     # --- the core fix: time from frames, vendor times ignored -----
48
49     def test_time_recomputed_from_frames_not_attributes(tmp_path):
50         # rally 1 spans source frames 60..120; the vendor's start/end attributes are
51         # deliberately wrong (the frame/900 bug). At 60 fps the truth is 1.0..2.0s.
52         xml = _write_cvat(tmp_path / "j.xml", [
53             _img_box(60, {"rally_number": "1.0", "start_time": "0.2", "end_time": "1.4",
54                 "shots_count": "6.0", "ending_reason": "winner",
55                 "score_before (A-B)": "0-0", "notes": "clean"}),
56             _img_box(90, {"rally_number": "1.0", "start_time": "0.2", "end_time": "1.4"}),
57             _img_box(120, {"rally_number": "1.0", "start_time": "0.2", "end_time": "1.4"}),
58         ])
59         res = convert(xml, fps=60.0)
60         assert len(res.records) == 1
61         r = res.records[0]
62         assert r.start_time == 1.0 and r.end_time == 2.0           # frames/fps, not 0.2/1.4
63         assert r.vendor_start_time == "0.2"                       # preserved for transparency
64         assert r.shots_count == 6
65         assert r.ending_reason == "winner"
66         assert r.score_before == "0-0"                            # "(A-B)" alias resolved
67         assert r.notes == "clean"
68
69
70     def test_step_and_fps_read_from_meta_and_arg(tmp_path):
71         xml = _write_cvat(tmp_path / "j.xml", [
72             _img_box(0, {"rally_number": "1", "shots_count": "3"}),
73             _img_box(30, {"rally_number": "1"}),
74         ], step=30)
75         res = convert(xml, fps=60.0)
76         assert res.step == 30

```

```

77         assert res.fps == 60.0
78
79
80     def test_fps_required_when_unavailable(tmp_path):
81         xml = _write_cvat(tmp_path / "j.xml", [_img_box(0, {"rally_number": "1"})])
82         try:
83             convert(xml) # no fps, no video_path
84         except ValueError as e:
85             assert "fps" in str(e).lower()
86         else:
87             raise AssertionError("expected ValueError when fps cannot be determined")
88
89
90     # --- track-mode parsing + outside filtering -----
91
92     def test_track_mode_and_outside_filtering(tmp_path):
93         xml = (
94             '<?xml version="1.0"?><annotations><meta><job>'
95             '<stop_frame>1000</stop_frame><frame_filter>step=30</frame_filter></job></meta>'
96             '<track id="0" label="rally">'
97             '<box frame="60" outside="0" xtl="1" ytl="1" xbr="2" ybr="2">'
98             '<attribute name="rally_number">1</attribute>'
99             '<attribute name="shots_count">4</attribute></box>'
100             '<box frame="120" outside="0" xtl="1" ytl="1" xbr="2" ybr="2">'
101             '<attribute name="rally_number">1</attribute></box>'
102             '<box frame="150" outside="1" xtl="1" ytl="1" xbr="2" ybr="2">' # terminator
103             '<attribute name="rally_number">1</attribute></box>'
104             '</track></annotations>'
105         )
106         p = tmp_path / "track.xml"
107         p.write_text(xml, encoding="utf-8")
108         meta, boxes = parse_cvat_boxes(str(p))
109         assert len(boxes) == 2 # the outside=1 terminator dropped
110         recs = build_records(boxes, fps=60.0)
111         assert recs[0].frame_max == 120 # not 150 (the terminator)
112         assert recs[0].end_time == 2.0
113
114
115     # --- issue detection: surfaced, not silently fixed -----
116
117     def test_overlap_is_a_blocker():
118         recs = [_rec(1, 1.0, 5.0), _rec(2, 4.0, 8.0)] # r2 starts before r1 ends
119         issues = detect_issues(recs)
120         overlaps = [i for i in issues if i.code == "overlap"]
121         assert len(overlaps) == 1 and overlaps[0].severity == "blocker"
122
123
124     def test_zero_duration_is_a_blocker():
125         recs = [_rec(1, 2.0, 2.0, fmin=120, fmax=120, n=1)]
126         issues = detect_issues(recs)
127         assert any(i.code == "zero_duration" and i.severity == "blocker" for i in issues)
128
129
130     def test_offvocab_ending_reason_flagged():
131         recs = [_rec(1, 1.0, 3.0, reason="into_net")]
132         issues = detect_issues(recs)
133         assert any(i.code == "ending_reason_vocab" for i in issues)
134         # the adopted forced/unforced terms are accepted, not flagged
135         ok = detect_issues([_rec(1, 1.0, 3.0, reason="forced_error")])
136         assert not any(i.code == "ending_reason_vocab" for i in ok)
137
138
139     def test_low_shot_density_flags_inflated_span():
140         recs = [_rec(1, 0.0, 16.0, shots=2)] # 2 shots over 16s -> 0.12/s
141         issues = detect_issues(recs)

```

```

142     assert any(i.code == "shot_density" for i in issues)
143
144
145 def test_outcome_vocab_flagged():
146     bad = detect_issues([_rec(1, 1.0, 3.0, outcome="long")])
147     assert any(i.code == "outcome_vocab" for i in bad)
148     # the agreed outcomes are accepted, not flagged
149     ok = detect_issues([_rec(1, 1.0, 3.0, outcome="out")])
150     assert not any(i.code == "outcome_vocab" for i in ok)
151
152
153 # --- end-to-end: emitted CSV imports via the existing import-local parser ---
154
155 def test_canonical_csv_round_trips_through_import_local(tmp_path):
156     xml = _write_cvat(tmp_path / "j.xml", [
157         _img_box(60, {"rally_number": "1", "start_time": "0.2", "end_time": "1.4",
158             "shots_count": "6", "ending_reason": "winner",
159             "last_shot_outcome": "in"}),
160         _img_box(120, {"rally_number": "1"}),
161         _img_box(180, {"rally_number": "2", "shots_count": "4",
162             "ending_reason": "forced_error", "last_shot_outcome": "out"}),
163         _img_box(240, {"rally_number": "2"}),
164     ])
165     res = convert(xml, fps=60.0)
166     assert res.n_blockers == 0
167     assert res.records[1].last_shot_outcome == "out" # parsed from the box attribute
168     out = tmp_path / "canonical.csv"
169     write_canonical_csv(res.records, str(out))
170
171     # Re-read with the actual import-local parser (self is unused -> None is safe).
172     items = SportsDataCLI._parse_local_annotations(None, "badminton", "vid1", str(out))
173     assert len(items) == 2
174     assert items[0].start_time == 1.0 and items[0].end_time == 2.0 # not 0.2/1.4
175     assert items[1].ending_reason == "forced_error"
176     assert items[1].last_shot_outcome == "out" # survives the round-trip into
the schema
177     assert items[1].start_time == 3.0 and items[1].end_time == 4.0
178
179
180 # --- validate-delivery: gap heuristic, CSV loading, report artifacts -----
181
182 def test_possible_missed_rally_gap_flagged():
183     # 20s of unlabeled time between two rallies -> flag (catches the 1:46-style miss)
184     recs = [_rec(1, 95.6, 98.0), _rec(2, 118.0, 119.6)]
185     issues = detect_issues(recs)
186     assert any(i.code == "possible_missed_rally" for i in issues)
187
188
189 def test_normal_gap_not_flagged():
190     recs = [_rec(1, 10.0, 14.0), _rec(2, 22.0, 26.0)] # 8s gap -> fine
191     issues = detect_issues(recs)
192     assert not any(i.code == "possible_missed_rally" for i in issues)
193
194
195 def test_load_canonical_csv_decimal_and_mmss(tmp_path):
196     p = tmp_path / "c.csv"
197     p.write_text(
198         "rally_number,start_time,end_time,start_mmss,end_mmss,shots_count,ending_reason,
last_shot_outcome\n"
199         "1,3.0,21.0,,15,unforced_error,out\n"
200         "2,,,0:29.000,0:32.000,2,winner,in\n", encoding="utf-8")
201     recs = load_canonical_csv(str(p))
202     assert len(recs) == 2
203     assert recs[0].start_time == 3.0 and recs[0].last_shot_outcome == "out"
204     assert recs[1].start_time == 29.0 and recs[1].end_time == 32.0 # mm:ss fallback

```

```

205
206
207 def test_load_canonical_csv_keeps_rally_starting_at_zero(tmp_path):
208     # Regression: a rally that legitimately starts at 0.0s must not be dropped.
209     # The old `start_time` or `start_mmss` fallback treated the falsy 0.0 as "missing"
210     # and skipped the row when there was no start_mmss column to fall back to.
211     p = tmp_path / "c.csv"
212     p.write_text(
213         "rally_number,start_time,end_time,shots_count,ending_reason,last_shot_outcome\n"
214         "1,0.0,4.0,8,winner,in\n"
215         "2,10.0,14.0,6,forced_error,out\n"
216         "3,20.0,24.0,5,unforced_error,out\n", encoding="utf-8")
217     recs = load_canonical_csv(str(p))
218     assert len(recs) == 3 # the zero-start row 1 is not lost
219     assert recs[0].start_time == 0.0 and recs[0].end_time == 4.0
220
221
222 def test_validation_artifacts(tmp_path):
223     recs = [_rec(1, 1.0, 5.0), _rec(2, 4.0, 8.0, reason="into_net")] # overlap + bad
vocab
224     issues = detect_issues(recs)
225     summary = build_summary(recs, issues, source="vid", fps=59.94)
226     assert summary["n_rallies"] == 2
227     assert summary["passed"] is False and summary["n_blockers"] >= 1
228     assert 2 in summary["flagged_rallies"]
229     md = render_validation_md(recs, issues, "vid")
230     assert "BLOCKED" in md and "## BLOCKER" in md
231     out = tmp_path / "review.csv"
232     write_review_csv(recs, issues, str(out))
233     rows = list(_csv.DictReader(open(out, encoding="utf-8")))
234     assert "VERIFIED(Y/N)" in rows[0] and "auto_flags" in rows[0]
235     assert any("overlap" in r["auto_flags"] for r in rows)
236

```

38. user (2026-06-10T06:03:12.534Z)

```

1     """normalize_annotations: the two-landmine guard for partial human annotator CSVs.
2
3     Pins the behavior that protects import-local: drop unusable rows (instead of aborting
4     the whole import) and renumber fractional/duplicate rally ids (instead of a silent PK
5     collision). The final test feeds the normalized output through the REAL import-local
6     parser to prove it ingests cleanly.
7     """
8     import csv as _csv
9
10    from src.cli.normalize_annotations import normalize_rows, normalize_file
11    from src.cli.curate import SportsDataCLI
12
13
14    def _rows(*dicts):
15        return [dict(d) for d in dicts]
16
17
18    def test_drops_blank_and_reversed_end_rows():
19        out, rep = normalize_rows(_rows(
20            {"rally_number": "1", "start_time": "10.0", "end_time": "15.0"}, # good
21            {"rally_number": "2", "start_time": "20.0", "end_time": ""}, # blank end
-> drop
22            {"rally_number": "3", "start_time": "30.0", "end_time": "28.0"}, # end<=start
-> drop
23            {"rally_number": "4", "start_time": "40.0", "end_time": "45.0"}, # good
24        ))
25        assert rep.rows_in == 4 and rep.rows_out == 2
26        assert len(rep.dropped) == 2
27        assert [r["start_time"] for r in out] == ["10.000", "40.000"]

```

```

28         # surviving rows are cleanly renumbered 1..2
29         assert [r["rally_number"] for r in out] == [1, 2]
30
31
32     def test_drops_negative_start_rows():
33         # BadmintonRally's validator rejects start<0 -> would abort the whole import.
34         out, rep = normalize_rows(_rows(
35             {"rally_number": "1", "start_time": "-5.0", "end_time": "10.0"}, # negative
start -> drop
36             {"rally_number": "2", "start_time": "20.0", "end_time": "25.0"}, # good
37         ))
38         assert rep.rows_out == 1
39         assert any(d["reason"] == "negative start_time" for d in rep.dropped)
40         assert out[0]["start_time"] == "20.000"
41
42
43     def test_renumbers_fractional_and_duplicate_ids_by_start():
44         out, rep = normalize_rows(_rows(
45             {"rally_number": "7", "start_time": "100.0", "end_time": "105.0"},
46             {"rally_number": "7.5", "start_time": "106.0", "end_time": "110.0"}, # inserted
correction
47             {"rally_number": "8", "start_time": "120.0", "end_time": "125.0"},
48         ))
49         assert rep.renumbered is True
50         assert rep.fractional_ids == ["7.5"]
51         # contiguous integer PKs in start order -> no (video_id, rally_number) collision
52         assert [r["rally_number"] for r in out] == [1, 2, 3]
53         assert [r["original_rally_number"] for r in out] == ["7", "7.5", "8"]
54
55
56     def test_duplicate_truncating_ids_are_separated():
57         # Two ids that BOTH int(float()) to 7 -> would collide on PK; must become 1,2.
58         out, rep = normalize_rows(_rows(
59             {"rally_number": "7.1", "start_time": "10.0", "end_time": "12.0"},
60             {"rally_number": "7.9", "start_time": "20.0", "end_time": "22.0"},
61         ))
62         assert [r["rally_number"] for r in out] == [1, 2]
63         assert rep.renumbered is True
64
65
66     def test_mmss_times_converted_and_counted():
67         out, rep = normalize_rows(_rows(
68             {"rally": "1", "start_mmss": "1:46.031", "end_mmss": "1:50.000"},
69         ))
70         assert out[0]["start_time"] == "106.031" and out[0]["end_time"] == "110.000"
71         assert rep.mmss_converted == 2
72
73
74     def test_pipe_delimited_reason_folds_to_primary_plus_note():
75         out, _ = normalize_rows(_rows(
76             {"rally_number": "1", "start_time": "1.0", "end_time": "2.0", "ending_reason":
"forced_error|other"},
77         ))
78         assert out[0]["ending_reason"] == "forced_error"
79         assert "other" in out[0]["notes"]
80
81
82     def test_clean_lean_input_is_passthrough():
83         out, rep = normalize_rows(_rows(
84             {"rally_number": "1", "start_time": "1.0", "end_time": "2.0", "ending_reason":
"winner", "sport": "badminton"},
85             {"rally_number": "2", "start_time": "3.0", "end_time": "5.0", "ending_reason":
"unforced_error", "sport": "badminton"},
86         ))
87         assert rep.rows_out == 2 and rep.renumbered is False and not rep.dropped

```

```

88         assert [r["sport"] for r in out] == ["badminton", "badminton"]
89
90
91     def test_normalized_output_ingests_through_import_local_parser(tmp_path):
92         """End-to-end: a dirty annotator CSV -> normalize -> the REAL parser yields N
93         BadmintonRally with contiguous int ids and no exception (the import-abort + PK
94         collision are both gone)."""
95         dirty = tmp_path / "dirty.rallies.csv"
96         with open(dirty, "w", newline="", encoding="utf-8") as f:
97             w = _csv.writer(f)
98             w.writerow(["rally_number", "start_time", "end_time", "ending_reason", "sport"])
99             w.writerow([1, 37.434, 72.194, "winner", "badminton"])
100             w.writerow([2, 79.936, "", "winner", "badminton"])          # blank end (would
abort)
101             w.writerow([2.5, 90.0, 95.0, "unforced_error", "badminton"]) # fractional
(would collide)
102             w.writerow([3, 130.193, 134.947, "winner", "badminton"])
103         clean = tmp_path / "clean.csv"
104         rep = normalize_file(str(dirty), str(clean))
105         assert rep.rows_out == 3 and len(rep.dropped) == 1
106
107         # self is unused inside _parse_local_annotations -> None is safe.
108         items = SportsDataCLI._parse_local_annotations(None, "badminton", "adarsh_avi",
str(clean))
109         assert len(items) == 3
110         assert [it.rally_number for it in items] == [1, 2, 3]    # contiguous, no collision
111         assert all(it.end_time > it.start_time for it in items)
112

```

39. user (2026-06-10T06:03:13.013Z)

```

1         """Integration tests for the import-local annotation parser on the newly-added
2         sports (pickleball, table tennis). _parse_local_annotations does not touch `self`,
3         so we can exercise it directly with a temp file and no DB side effects."""
4         import json
5         import csv as _csv
6
7         from src.cli.curate import SportsDataCLI
8         from src.core.schemas.pickleball import PickleballRally
9         from src.core.schemas.tabletennis import TableTennisRally
10
11
12     def _parse(sport, video_id, path):
13         # self is unused inside _parse_local_annotations, so None is safe here.
14         return SportsDataCLI._parse_local_annotations(None, sport, video_id, str(path))
15
16
17     def test_parse_pickleball_csv(tmp_path):
18         p = tmp_path / "pb.csv"
19         with open(p, "w", newline="", encoding="utf-8") as f:
20             w = _csv.writer(f)
21             w.writerow(["rally_number", "start_time", "end_time", "shots_count",
"ending_reason", "score_before", "score_after", "notes"])
22             w.writerow([1, 12.50, 25.00, 6, "winner", "0-0", "1-0", ""])
23             w.writerow([2, 30.00, 31.20, 1, "service_fault", "1-0", "1-1", "serve into
net"])
24
25         items = _parse("pickleball", "pb_csv", p)
26         assert len(items) == 2
27         assert all(isinstance(i, PickleballRally) for i in items)
28         assert items[0].shots_count == 6
29         assert items[1].ending_reason == "service_fault"
30         # Unknown columns (notes) are ignored, not fatal.
31         assert items[0].start_time == 12.5
32

```



```

33
34     def test_parse_tabletennis_json(tmp_path):
35         p = tmp_path / "tt.json"
36         data = [
37             {"rally_number": 1, "start_time": 5.0, "end_time": 9.0, "shots_count": 7,
"ending_reason": "winner"},
38             {"rally_number": 2, "start_time": 12.0, "end_time": 13.0, "shots_count": 1,
"ending_reason": "let"},
39         ]
40         p.write_text(json.dumps(data), encoding="utf-8")
41
42         items = _parse("tabletennis", "tt_json", p)
43         assert len(items) == 2
44         assert all(isinstance(i, TableTennisRally) for i in items)
45         assert items[1].ending_reason == "let"
46         assert items[0].video_id == "tt_json"
47

```

40. user (2026-06-10T06:03:13.684Z)

```

1     import os
2     import pytest
3     from src.core.db import SportsDatabase
4     from src.core.schemas.base import VideoMetadata, SportType, LicenseType, CurationStatus
5     from src.core.schemas.badminton import BadmintonRally
6     from src.core.schemas.pickleball import PickleballRally
7     from src.core.schemas.tabletennis import TableTennisRally
8     from src.core.schemas.tennis import TennisPoint
9     from src.core.schemas.cricket import CricketDelivery
10
11     @pytest.fixture
12     def temp_db(tmp_path):
13         db_file = tmp_path / "test_sports.db"
14         return SportsDatabase(str(db_file))
15
16     def test_insert_and_get_video(temp_db):
17         video = VideoMetadata(
18             video_id="test_match_1",
19             sport=SportType.BADMINTON,
20             video_url="http://youtube.com/watch?v=123",
21             local_path=None,
22             license_type=LicenseType.MIT,
23             license_url=None,
24             is_commercial=True,
25             status=CurationStatus.STAGING,
26             match_name="Axelsen vs Momota",
27             fps=30.0,
28             resolution="1080p"
29         )
30
31         assert temp_db.insert_video(video) is True
32
33         fetched = temp_db.get_video("test_match_1")
34         assert fetched is not None
35         assert fetched.match_name == "Axelsen vs Momota"
36         assert fetched.is_commercial is True
37         assert fetched.status == CurationStatus.STAGING
38
39     def test_badminton_rally_insertion(temp_db):
40         video = VideoMetadata(
41             video_id="match_2",
42             sport=SportType.BADMINTON,
43             license_type=LicenseType.MIT,
44             is_commercial=True,
45             match_name="Match 2"

```

```

46         )
47         temp_db.insert_video(video)
48
49         rallies = [
50             BadmintonRally(video_id="match_2", rally_number=1, start_time=10.0,
51                             end_time=25.0, score_before="0-0", score_after="1-0", shots_count=5),
52             BadmintonRally(video_id="match_2", rally_number=2, start_time=30.0,
53                             end_time=45.0, score_before="1-0", score_after="1-1", shots_count=12)
54         ]
55         inserted = temp_db.insert_badminton_rallies(rallies)
56         assert inserted == 2
57
58         fetched_rallies = temp_db.get_badminton_rallies("match_2")
59         assert len(fetched_rallies) == 2
60         assert fetched_rallies[0].rally_number == 1
61         assert fetched_rallies[1].shots_count == 12
62
63     def test_reinsert_video_preserves_child_rallies(temp_db):
64         """Re-inserting an existing video must update it in place, not wipe its rallies (bug
65         3)."""
66         video = VideoMetadata(
67             video_id="match_upsert",
68             sport=SportType.BADMINTON,
69             license_type=LicenseType.MIT,
70             is_commercial=True,
71             match_name="Original Name",
72         )
73         temp_db.insert_video(video)
74         temp_db.insert_badminton_rallies([
75             BadmintonRally(video_id="match_upsert", rally_number=1, start_time=1.0,
76                             end_time=5.0),
77         ])
78
79         # Re-insert the same video_id with a changed field.
80         video.match_name = "Updated Name"
81         temp_db.insert_video(video)
82
83         fetched = temp_db.get_video("match_upsert")

```

41. assistant (2026-06-10T06:03:21.188Z)

Now let me empirically verify the landmine behaviors with a temp DB (safe ? temp dir only).

42. user (2026-06-10T06:03:32.193Z)

```

safe_int('7.5') = 7
CASE1: RAISED IntegrityError UNIQUE constraint failed: badminton_rallies.video_id,
badminton_rallies.rally_number
DB rows after exception: 0
CASE2 before: 40
CASE2 after re-import of 1-row file: 1 (silent replace, no warning)
CASE3: video row silently overwritten -> sport: tennis | match: totally different match |
badminton rallies still attached: 1

```

43. assistant (2026-06-10T06:03:47.400Z)

Confirmed. Let me verify the rollback behavior on a failed import over existing golden rows, and corrupt-XML handling.

44. user (2026-06-10T06:03:58.500Z)

```

before failed import: 3
raised: IntegrityError

```

after failed import (rollback? old rows preserved?): 3
corrupt XML -> ParseError | caught by convert_cvat's except (FileNotFoundError, ValueError)?
False
empty CVAT export -> records: 0 issues: 0 (no error, silent empty)
space-named attrs matched? None

45. user (2026-06-10T06:05:04.430Z)

Name	Stmts	Miss	Cover
src\cli\curate.py	567	180	68%
src\cli\normalize_annotations.py	133	35	74%
src\core\config.py	15	0	100%
src\core\db.py	191	7	96%
src\core\downloader.py	82	9	89%
src\core\parsing.py	35	2	94%
src\core\schemas\badminton.py	19	0	100%
src\core\schemas\base.py	36	0	100%
src\core\schemas\cricket.py	20	0	100%
src\core\schemas\pickleball.py	19	0	100%
src\core\schemas\tabletennis.py	19	0	100%
src\core\schemas\tennis.py	18	0	100%
src\core\validator.py	88	5	94%
src\parsers\badminton\coachai.py	62	4	94%
src\parsers\base_parser.py	7	1	86%
src\parsers\cvat__init__.py	0	0	100%
src\parsers\cvat\rally_cvat.py	251	37	85%
src\reporting__init__.py	24	2	92%
src\reporting\adapter.py	76	3	96%
TOTAL	1662	285	83%

124 passed in 5.13s

099cf67 Manual-ingestion: normalize-annotations guard + positioning doc (#13)
0e4ccc7 Merge pull request #12 from avidullu/feat/collector-report-refinements
b8d40c8 feat(reporting): delivery-report debuggability refinements + faq-resolution golden test
94fec51 Merge pull request #11 from avidullu/feat/observability-reports
038424d feat(reporting): per-video observability report (collector)
99dfa22 Merge pull request #10 from avidullu/fix/canonical-csv-zero-start
f7cab11 fix(parser): keep a rally that starts at 0.0s in load_canonical_csv
d68c5da Merge pull request #9 from avidullu/feat/validate-delivery
eval_export
mock_local_badminton_anno.json
roora_badminton_canonical.csv
roora_badminton_issues.txt
sports_metadata.db
staging_coachai_match.csv
temp
validation
videos
badminton
manifest.json

46. user (2026-06-10T06:05:26.555Z)

src\cli\curate.py	567	180	68%	35, 166, 185, 209, 263, 287-288, 329-334, 337-339, 361, 371-373, 377, 380-381, 386-387, 391-392, 398-401, 409, 425, 457, 469-471, 479-480, 484-485, 533, 535, 549-582, 595-656, 687-688, 690-691, 702-709, 727-738, 761, 773-774, 783-814, 817-818, 825-827, 835, 841, 948, 950, 960
src\cli\normalize_annotations.py	133	35	74%	56-74, 83, 88, 91, 93-94, 103, 106, 132-134, 207-214, 218
TOTAL	700	215	69%	

124 passed in 5.82s

```
{  
  "commercial_only": true,
```

```

"statuses": [
  "curated",
  "staging"
],
"total_videos": 24,
"total_segments": 1915,
"eval_sets": [
  {
    "video_id": "my_custom_local_match",
    "sport": "badminton",
    "csv": "badminton/my_custom_local_match.csv",
    "local_path": "data/videos/my_custom_match.mp4",
    "video_url": null,
    "match_name": "My Custom Local Match",
    "fps": 30.0,
    "resolution": "1080p",
    "license_type": "Proprietary",
    "is_commercial": true,
    "status": "curated",
    "num_segments": 2
  },
  {
    "video_id": "shuttleaset_21",
    "sport": "badminton",
    "csv": "badminton/shuttleaset_21.csv",
    "local_path": "./data/videos\\shuttleaset_21.mp4",
    "video_url": "https://www.youtube.com/watch?v=gloiZ_gTJaE",
    "match_name": "An Se Young Ratchanok Intanon YONEX Thailand Open 2021 QuarterFinals",
    "fps": 30.0,
    "resolution": "1080p",
    "license_type": "MIT",
    "is_commercial": true,
    "status": "curated",
    "num_segments": 75
  },
  {
    "video_id": "shuttleaset_1",
    "sport": "badminton",
    "csv": "badminton/shuttlemy_custom_local_match.csv
shuttleaset_1.csv
shuttleaset_10.csv
shuttleaset_11.csv
shuttleaset_12.csv
shuttleaset_13.csv
shuttleaset_14.csv
shuttleaset_15.csv
shuttleaset_16.csv
shuttleaset_17.csv
shuttleaset_18.csv
shuttleaset_19.csv
shuttleaset_2.csv
shuttleaset_20.csv
shuttleaset_21.csv
shuttleaset_22.csv
shuttleaset_23.csv
shuttleaset_24.csv
shuttleaset_25.csv
shuttleaset_3.csv
shuttleaset_4.csv
shuttleaset_6.csv
shuttleaset_8.csv
shuttleaset_9.csv

```

```

data/mock_local_badminton_anno.json
data/roora_badminton_canonical.csv
data/roora_badminton_issues.txt
data/sports_metadata.db
data/staging_coachai_match.csv
data/validation/pilot_badminton_01_report.json
data/validation/pilot_badminton_01_report.md
data/validation/pilot_badminton_01_review.csv
---untracked-data---
pilot_badminton_01_report.json
pilot_badminton_01_report.md
pilot_badminton_01_review.csv
---canonical head---
rally_number,start_mmss,end_mmss,start_time,end_time,shots_count,ending_reason,last_shot_outcome
,score_before,score_after,notes
1,0:03.003,0:21.021,3.003,21.021,15,other,,,,net
2,0:02.503,0:32.533,2.503,32.533,2,forced_error,,,,

```

48. user (2026-06-10T06:05:50.021Z)

```

Found 2 files
tests\test_wasb_cache.py
tests\test_eval_batch.py

```

49. user (2026-06-10T06:06:04.431Z)

```

import os

from backend.eval import batch

```

---- path resolution -----

```

def test_resolve_relative_windows_path():
    got = batch.resolve_video_path("./data/videos\\shuttleaset_1.mp4", "/collector")
    assert got == os.path.normpath("/collector/data/videos/shuttleaset_1.mp4")

def test_resolve_absolute_path_kept():
    got = batch.resolve_video_path("C:/vids/x.mp4", "/collector")
    assert got == os.path.normpath("C:/vids/x.mp4")

def test_resolve_none():
    assert batch.resolve_video_path(None, "/collector") is None
    assert batch.resolve_video_path("", "/collector") is None

```

---- stage selection -----

```

def test_auto_stage_final_only_for_motion():
    assert batch.default_stages_for("motion", "auto") == ["final"]
    assert batch.default_stages_for("gemini", "auto") == ["final"]

def test_auto_stage_both_for_two_stage_detector():
    assert batch.default_stages_for("yolo_hybrid", "auto") == ["candidates", "final"]

def test_explicit_stage_overrides_auto():
    assert batch.default_stages_for("motion", "both") == ["candidates", "final"]
    assert batch.default_stages_for("yolo_hybrid", "final") == ["final"]

```

---- aggregation -----

```
def _vid(tp, fp, fn, miou, f1):
    return {"stages": {"final": {"score": {
        "true_positives": tp, "false_positives": fp, "false_negatives": fn,
        "mean_iou": miou, "f1": f1,
    }}}}
```

```
def test_aggregate_micro_and_macro():
    per_video = [
        _vid(tp=8, fp=2, fn=2, miou=0.9, f1=0.8), # P=.8 R=.8 F1=.8
        _vid(tp=6, fp=0, fn=4, miou=0.8, f1=0.75), # P=1 R=.6
    ]
    agg = batch.aggregate(per_video, ["final"])["final"]
    assert agg["tp"] == 14 and agg["fp"] == 2 and agg["fn"] == 6
    assert agg["videos"] == 2
    # micro precision = 14/16, recall = 14/20
    assert round(agg["micro_precision"], 3) == round(14 / 16, 3)
    assert round(agg["micro_recall"], 3) == round(14 / 20, 3)
    # TP-weighted mean IoU = (0.9*8 + 0.8*6)/14
    assert round(agg["mean_iou"], 4) == round((0.9 * 8 + 0.8 * 6) / 14, 4)
```

50. user (2026-06-10T06:06:04.750Z)

```
1      """Best-quality video acquisition.
2
3      A single place that knows how to fetch the *highest resolution available* for a
4      source URL and report what was actually obtained. Both the ShuttleSet crawler
5      (initial ingestion) and the `fetch` CLI command (in-place upgrade of already-
6      ingested videos) go through here, so the resolution policy lives in exactly one
7      spot.
8
9      Design notes:
10     - `bv*+ba/b` selects best video + best audio, falling back to the best
11       pre-muxed stream. Combined with `-S res,fps,...`, yt-dlp sorts candidate
12       formats by resolution (then fps) descending and picks the top ? i.e. *always*
13       the best resolution available, regardless of codec/container quirks.
14     - `--merge-output-format mp4` guarantees an MP4 container, which the downstream
15       indexer requires.
16     - We never trust the requested resolution; we ffprobe the resulting file and
17       record the *actual* width/height/fps as the source of truth.
18     """
19
20     import os
21     import shutil
22     import subprocess
23     from dataclasses import dataclass
24     from typing import List, Optional, Tuple
25
26     # Best video+audio, prefer mp4/m4a, fall back to best single stream.
27     _FORMAT = "bv*+ba/b"
28     # Sort candidates so the highest resolution (then fps) wins; prefer mp4/m4a on ties.
29     _SORT = "res,fps,ext:mp4:m4a"
30
31
32     def height_to_label(height: Optional[int]) -> Optional[str]:
33         """Map a pixel height to a conventional resolution label (e.g. 1080 -> '1080p').
34
35         Snaps to the nearest standard tier at or below the actual height so a 1072px
36         crop still reads as '1080p' rather than an odd number.
37         """
38         if not height or height <= 0:
39             return None
```

```

40     # Allow a 5% tolerance below each tier so a lightly letterboxed/cropped frame
41     # (e.g. 1072px) still reads as the intended tier ('1080p') rather than dropping.
42     for tier in (2160, 1440, 1080, 720, 480, 360, 240, 144):
43         if height >= tier * 0.95:
44             return f"{tier}p"
45     return f"{height}p"
46
47
48 def build_ytdlp_command(url: str, dest_path: str, max_height: Optional[int] = None,
49                        cookies_from_browser: Optional[str] = None,
50                        cookies_file: Optional[str] = None,
51                        remote_components: Optional[str] = "ejs:github",
52                        player_client: Optional[str] = "default") -> List[str]:
53     """Build the yt-dlp argv for a best-quality download.
54
55     `max_height`, when set, caps resolution (e.g. avoid pulling 4K); left unset
56     it fetches the absolute best available. Pure/argv-only so it is unit-testable
57     without invoking the network.
58
59     `player_client` selects YouTube's player client (default `default` = yt-dlp's
60     maintained client set, which serves real downloadable 1080p DASH with auth
61     cookies). Avoid `web`/`web_safari`: YouTube forces those into SABR streaming,
62     which strips download URLs and silently drops to the lowest stream (the 144p
63     bug). Harmless for non-YouTube URLs (the extractor-arg is ignored). Set to
64     None to omit the arg entirely.
65
66     `remote_components` enables yt-dlp's EJS challenge solver (default
67     `ejs:github`), the JS-runtime (deno) path for solving YouTube's n-challenge.
68     Set to None to disable (e.g. for non-YouTube sources or offline tests).
69     """
70     fmt = _FORMAT
71     if max_height:
72         # Cap each stream selector; keep the pre-muxed fallback capped too.
73         fmt = f"bv*[height<={max_height}]+ba/b[height<={max_height}]/b"
74     cmd = [
75         "yt-dlp",
76         "-f", fmt,
77         "-S", _SORT,
78         "--merge-output-format", "mp4",
79         "--no-playlist",
80         "--no-progress",
81         "-o", dest_path,
82     ]
83     if player_client:
84         cmd += ["--extractor-args", f"youtube:player_client={player_client}"]
85     if remote_components:
86         cmd += ["--remote-components", remote_components]
87     if cookies_from_browser:
88         cmd += ["--cookies-from-browser", cookies_from_browser]
89     if cookies_file:
90         cmd += ["--cookies", cookies_file]

```

51. user (2026-06-10T06:06:23.089Z)

CVAT rally conversion report ? annotations.xml

```

fps used      : 59.940
sampling step : 30 frames (~0.50s)
rallies parsed : 24
issues       : 5 (1 blocker)

```

[BLOCKER]

- (overlap) Rally 1 (ends 21.02s) overlaps Rally 2 (starts 2.50s) by 18.52s ? rallies must be disjoint.

[HIGH]

- (coarse_sampling) Frames sampled every 30 (~0.50s at 59.94fps); rally boundaries carry +/-0.50s uncertainty, exceeding the +/-0.3s spec.

[MEDIUM]

- (shot_density) Rally 2: 2 shots over 30.0s (0.07/s) is implausibly low ? the span may be inflated by a stray box.
- (shot_density) Rally 7: 2 shots over 16.0s (0.12/s) is implausibly low ? the span may be inflated by a stray box.
- (shot_density) Rally 16: 2 shots over 17.5s (0.11/s) is implausibly low ? the span may be inflated by a stray box.

===

Delivery validation ? pilot_badminton_01

- rallies: **24**
- issues: **4** (blocker 1, high 0, medium 3, low 0)
- structurally-clean rallies: **21/24** (87%)
- gate: **BLOCKED**

BLOCKER

- (overlap) Rally 1 (ends 21.02s) overlaps Rally 2 (starts 2.50s) by 18.52s ? rallies must be disjoint.

MEDIUM

- (shot_density) Rally 2: 2 shots over 30.0s (0.07/s) is implausibly low ? the span may be inflated by a stray box.

52. assistant (2026-06-10T06:07:24.299Z)

I have everything needed. Compiling the audit findings.

53. user (2026-06-10T06:09:13.242Z)

Structured output provided successfully

54. assistant (2026-06-10T06:09:27.951Z)

Audit complete. Structured output delivered: 15 findings (1 critical ? licensing conflation; 4 high ? silent destructive replace, empty-delivery PASS gate, missing SoR provenance/maturity, rally_number truncation crash), 8 strengths, all with file/line evidence, test counts (124 passed, 83% coverage measured) and empirically-verified landmine behaviors.