

claude-windows / 45ea3a91-654d-4972-9cd3-65f35db54caf

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\45ea3a91-654d-4972-9cd3-65f35db54caf\subagents\workflows\wf_2a64bc88-2a8\agent-a489ea466c8fd021f.jsonl`  
- SHA-256: `a5f13531f5274cad057a2a39c20b1781d1342f76412311ae1e990f76598349a5`  
- Source modified: `2026-06-10T06:12:46+00:00`  
- Imported at: `2026-07-05T16:48:27+00:00`  
- project: `wf_2a64bc88-2a8`  
- session_id: `45ea3a91-654d-4972-9cd3-65f35db54caf`
```

Transcript

1. user (2026-06-10T05:59:28.041Z)

CONTEXT ? two sibling repos on a Windows machine:

- INDEXER: C:/Users/avidu/Projects/badminton-highlight-indexer ? local AI badminton(->racket-sports) highlight indexer + rally detector. FastAPI + SQLite + ffmpeg + GPU CV via WSL; Gemini = paid cloud AI. docs/ tree is rich. Current branch docs/next-steps-multisport (master = main branch).

- COLLECTOR: C:/Users/avidu/Projects/sports-data-collector ? golden-label acquisition/curation/ingestion CLI (system-of-record candidate for training corpus).

Project state (June 2026): scaffolding complete; owned-model roadmap agreed (docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md) but NO platform/owned-model implementation started; next committed platform slice = async job model. Licensing guardrail: MIT/Apache-2.0/BSD only for code+data+weights; paid LLMs are inference-only, never training-data sources. ENVIRONMENT RULES: this dev environment has NO GPU/torch/cv2 ? do NOT attempt to run the video pipeline or import torch/cv2 modules. Test suites are offline/fully-mocked and SAFE to run. You are READ-ONLY: never edit/create/delete repo files; running tests and git read commands is allowed.

QUALITY BAR: cite file paths (and line numbers where possible) for every finding. Report what IS in the code/docs, not what you assume. Be skeptical and concrete. Your final structured output is consumed programmatically by an orchestrator ? return raw data, not prose for a human.

TASK: deep ENGINEERING audit of the indexer's runtime pipeline code. Read docs/CODE_MAP.md FIRST (dense navigation map with known gotchas ? do not just re-report what it already flags; confirm and go DEEPER). Then deep-read: backend/main.py, backend/config/ (models.py, loader.py), backend/pipeline/segmenters/ (base + all 5), backend/pipeline/detectors/ (base, tracknet_runner, wasb_runner, wasb_infer, rally_gate), backend/pipeline/validators/ (all), backend/pipeline/stitcher/compiler.py, backend/pipeline/audio/hit_detector.py, backend/infrastructure/database.py, backend/providers/, backend/sports/, backend/annotations/, backend/reporting/ (incl. spec.yaml), backend/utils/.

ASSESS: (1) architecture quality ? registry/ABC plugin design, layering, seams for the planned platform work (StorageBackend/ComputeBackend, async jobs); (2) correctness risks ? real bugs, copy-paste drift between twin segmenters, contract mismatches; (3) code duplication and dead code; (4) API design of the FastAPI surface (validation, error responses, blocking behavior); (5) DB schema + migration ladder quality; (6) config system (shallow merge, extra=allow, .get shim) ? is it a time bomb or fine; (7) readiness for the NEXT committed slice (async job model) ? what will hurt.

Severity-rate findings honestly: 'critical' = would corrupt data / break prod runs; 'high' = real bug or design flaw with concrete consequence; 'medium' = drift/fragility; 'low'/'info' = polish. NEW findings (not already in CODE_MAP gotchas) are worth the most; if you confirm a CODE_MAP gotcha matters more than documented, say so explicitly.

2. user (2026-06-10T05:59:34.068Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-255 of 330 total (27423 tokens, cap 25000). Call Read with offset=256 limit=255 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-

reminder>

```
1      # CODE_MAP.md
2
3      **Purpose:** cold-start navigation map for the badminton-highlight-indexer ? a self-
hosted video ? highlight-reel pipeline (CV/AI rally detection ? SQLite ? ffmpeg stitch) with a
segmenter-agnostic eval/calibration stack, typed config, and per-video observability reporting.
4
5      **Regenerate:** rerun the `build-code-map` workflow (parallel subsystem readers ? this
doc). Re-run after any registry/contract/WSL-interface change, or after a directory restructure
(this revision follows #40).
6
7      **LEGEND**
8      - `@path` = file (always repo-relative)
9      - `::sym` = symbol (class/func/const) in the nearest preceding `@path`
10     - `->` = dataflow / call / "produces"
11     - `$` = data contract (canonical shape)
12     - `?` = gotcha / footgun
13     - `?` = registry / plugin extension point
14     - `WSL` = runs inside WSL2 conda env, NOT the Windows Python process
15
16     ---
17
18     ## 1. PIPELINE
19
20     ### 1A. Runtime: video file -> highlight reel
21     ```
22     typed config load (built-in defaults -> config.json overrides; absent/parse-fail -> all-
defaults, never fatal)
23                                     @backend/config/loader.py::load_config ->
@backend/config/models.py::AppConfig
24     -> video_id = MD5(whole file)          @backend/utils/hashing.py::compute_video_id
(single shared helper, 64KB chunks; imported by main + all run_*.py ? no more per-entrypoint
copies)
25     -> validate
@backend/pipeline/validators/base.py::registry.run_all_with_context(video_path, global_config)
26     FormatValidator -> ctx['ffprobe_meta'] -> ResolutionFPSValidator -> PTSContinuity
-> SceneCut -> Blur -> Relevance (? ACTUAL registration/exec order ? see
@backend/pipeline/validators/__init__.py)
27     [validation FAIL -> persist status="failed", return early; report STILL emitted
in finally]
28     -> db.add_video(video_id, filepath, status, [r.dict()])
@backend/infrastructure/database.py
29     -> seg_name = req.segmenter or global_config.indexing.default_segmenter (fallback
'motion')
30     -> seg = segmenter_registry.get(seg_name)(db, global_config) ?
@backend/pipeline/segmenters/base.py::segmenter_registry (400 + available list if missing)
31     -> segments = seg.process_video(video_id, video_path, player_pool?, max_frames?)
32     [STOP POINT: segments-only ? predictions now in DB; eval/regression operate here
without stitching]
33     -> finally: emit_indexer_report(...) @backend/reporting/__init__.py (? ALWAYS runs
? even on validation failure / exception; NEVER raises; grafts response["reports"])
34     -> select segment ids -> [(start,end)] (+ stitching padding)
35     -> VideoCompiler.compile_highlights(filepath, time_pairs, out_name)
@backend/pipeline/stitcher/compiler.py
36     per-clip ffmpeg re-encode (libx264/superfast/crf22) -> concat -c copy ->
output/<name>.mp4
37     ```
38     Hybrid segmenter internal 4-phase (yolo_hybrid / tracknet_hybrid / wasb_hybrid ? same
shape):
39     ```
40     P1 chunk video (yolo only; tracknet/wasb = whole-video single pass, chunk_index=0)
41     -> P2 candidate "action windows" from motion/trajectory [STOP: candidates persisted
via db.add_candidate_segment]
42     -> P3 pad clip (ai_handoff_padding) -> ThreadPoolExecutor ->
```

```

provider.analyze_video(clip, prompt) ? provider_registry
43     -> P4 db.add_play_segment + db.link_candidate_to_segment
44     ...
45     Pure-AI path `gemini`: no P2; chunk-or-single -> provider -> heavy validate
(clamp/dedup) -> add_play_segment.
46     Legacy `motion`: pixel-diff -> segments + **fabricated** metadata/events (? not ground
truth).
47
48     ### 1B. WASB shuttle substrate (P2 of wasb_hybrid)
49     ...
50     WasbHybridSegmenter.process_video
@backend/pipeline/segmenters/wasb_hybrid.py
51     -> WasbRunner.healthcheck() (gate; not ok -> return [])
@backend/pipeline/detectors/wasb_runner.py
52     -> WasbRunner.run_predict(video, out_dir)
53     stage video + wasb_infer.py into WSL fs -> `wsl bash -c 'source conda.sh && python
wasb_infer.py ...'`
54     -> @backend/pipeline/detectors/wasb_infer.py::run (WSL): sliding windows -> GPU
detector pass (CHECKPOINTED det_raw.jsonl) -> CPU tracker (always full re-run) -> trajectory CSV
55     -> TrackNetRunner.parse_trajectory_csv(csv)
@backend/pipeline/detectors/tracknet_runner.py (shared, re-exported by WasbRunner)
56     -> TrackNetRunner.trajectory_to_action_windows(points, fps, frame_width, chunk_start,
velocity_thresh, merge_gap, min_window_duration, chunk_index) [the model-agnostic WINDOWING
BRAIN]
57     ...
58     (`tracknet_hybrid` identical, swapping WasbRunner->TrackNetRunner; `_probe_fps_width`
fallback (30.0, 1280.0).)
59
60     ### 1C. Calibration / eval flow (NO pipeline run unless told)
61     ...
62     GT CSV -> annotations.load_annotations / calibrate_wasb.load_golden_gts ->
List[Interval]
63     DB predictions -> harness.get_predictions(db, video_id, stage) -> List[Interval] stage
? {candidates, final}
64     -> metrics.match_intervals (greedy temporal-IoU) -> metrics.score -> Score
65     -> harness.evaluate -> EvalReport -> report_to_dict
66     -> batch.aggregate (corpus micro/macro) @backend/eval/batch.py
67     -> regression.regress(_with_provider) vs baseline @backend/eval/regression.py
[CI gate]
68     Calibration: calibrate_wasb.run -> calibration.{select_best, improvement_report,
cross_validate} @backend/eval/calibration.py
69     Distill Gemini->local: calibrate_local (thresholds) OR distill_local (learned
classifier) @backend/eval/{calibrate_local,distill_local}.py
70     Learned segmenter (offline): training.build_training_set(X,y) ->
classifier.LogisticRegression.fit/save
71     Gemini refinement driver (tuning, standalone):
gemini_refine.{refine_window,full_video_rallies} @backend/eval/gemini_refine.py
72     Audio probe (diagnostic only, NOT in live pipeline): audio/e1_probe.run
@backend/eval/audio/e1_probe.py
73     ...
74     Entrypoints: `@run_eval.py` (single video score), `@run_phase3_eval.py` (manifest batch,
can segment), `@run_regression.py` (baseline gate, exit 1=regression),
`@backend/eval/calibrate_wasb.py` + `@backend/eval/export_wasb_segments.py` (WASB tuning
drivers).
75
76     ---
77
78     ## 2. SUBSYSTEMS
79
80     ### 2.0 Orchestration & config
81     - `@backend/main.py` ? Thin CLI + FastAPI bootstrap (orchestrator; modes via FLAGS).
Static UI mounts `/static` from `backend/api/static`. Logging is initialized with `basicConfig`
(INFO, timestamped format); backend modules use `logger = logging.getLogger(__name__)` (replaces
stray prints in hot paths like motion, stitcher, wasb_infer, config loader). User-facing CLI
summaries may still use `print` for direct output.

```

```

82     - `::app` FastAPI; `::db_instance`, `::global_config` (`Optional[AppConfig]`) module
globals set ONLY in `main()`. ? importing `app` for ASGI without `main()` -> both `None` ->
every endpoint NPEs.
83     - `::process_video` `@POST /api/process` -> validate -> add_video -> segment ->
`finally:` always `emit_indexer_report` (never raises). `::compile_highlights` `@POST
/api/compile`. `::query_play_segments` `@POST /api/query`. `::get_config` `@GET /api/config`.
`::run_cli_diagnose_clip` `--debug-clip`: lazy `import cv2`, samples  $\pm 3s$  vs
`indexing.motion_threshold` (dual model/dict read).
84     - `::load_config` ? ? legacy thin wrapper returning a dict
(`load_app_config(path).model_dump(...)`); kept for transition ? new code imports
`load_app_config`/`AppConfig` directly. `::ProcessRequest`/`::QueryRequest`/`::CompileRequest`
pydantic bodies.
85     - `main()` sets `global_config.output.output_dir = args.output_dir` (typed field on
`OutputConfig`); `db_path = os.path.join(global_config.output.output_dir, output.db_name)`.
`compile_highlights` reads the SAME `global_config.output.output_dir`, so `--output-dir` is
honored end-to-end (the old write-only `_runtime_overrides` hack is gone).
86     - ? `--segmenter` argparse `choices` frozen at build time from
`segmenter_registry.list_available()` ? a lazily-registered segmenter wouldn't appear. CLI
default_segmenter fallback `motion`.
87     - `@run_process_and_stitch.py` ? one-shot demo; ? hardcoded paths, interactive `input()`
(blocks automation), skips validation. Config via `backend.config.load_config` (typed);
segmenter, padding, and db path all read from the model (no literal fallbacks).
88     - `@run_eval.py::main` ? score one video's DB preds vs GT CSV.
`--stage{candidates,final,both}`. Exit 2=arg/IO. ? pure scorer; errors if video never processed.
`::default_db_path` resolves via `backend.config.load_config`
(`output.output_dir`/`output.db_name`). `get_file_md5` is an alias for `compute_video_id`.
89     - `@run_phase3_eval.py::main` ? walk manifest, segment-if-needed, score, aggregate.
`load_dotenv()` at import. ? DB key = file MD5, NOT manifest's `video_id` (kept as label);
config via `backend.config.load_config`; `segmenter_cls(db, cfg)` receives the typed `AppConfig`
(same object the live pipeline feeds segmenters).
90     - `@run_regression.py::main` ? `regress_with_provider` vs baseline. Exit codes semantic:
0=ok, 1=REGRESSION (CI gate), 2=missing DB. ? `--update-baseline` runs AFTER compare
(snapshots even a regressing run); `::default_db_path` resolves via
`backend.config.load_config`.
91
92     - `@backend/config/__init__.py` ? ? canonical config import surface. Re-exports
`load_config`, `save_default_config` (from `.loader`) + `AppConfig`, `Config`, `IndexingConfig`,
`ValidationConfig` (from `models`); all = exactly those 6. Use `from backend.config
import load_config, AppConfig`.
93     - `@backend/config/models.py` ? Pydantic v2 models; single source of truth for
defaults.
94     - `::AppConfig` ? root. Scalars `sport='badminton'`, `ai_provider='gemini'`,
`ai_model='gemini-2.5-flash'`; sections `validation`, `indexing`, `output`, `stitching`
(`default_factory`). `model_config = {extra:'allow', validate_assignment:True}`. ?
`extra='allow'` -> unknown/typo'd config.json keys silently kept (legacy tolerance).
95     - ? `::AppConfig.get` ? legacy dict-compat shim: returns nested sections as plain
dicts (`model_dump(mode='json')`, NOT the model) so chained
`config.get("indexing",{}).get(...)` keeps working; falls back to searching a leaf key in any
section (sections discovered dynamically from `model_fields`). ? "bit every validator on real
runs" per docstring ? returning the model breaks `.get`. `::AppConfig.__getitem__` raises
`KeyError` for unknown top-level keys else delegates to `.get`.
96     - `::IndexingConfig` ? segmenter/detector knobs (`motion_threshold=0.08`,
`min_segment_duration=3.0`, `dead_time_merge_gap=2.0`, `chunk_duration=90.0`,
`chunk_overlap=10.0`, `yolo_velocity_threshold=0.15`, `yolo_frame_skip=3`,
`ai_handoff_padding=2.0`, `ai_max_workers=5`, `log/store_yolo_candidates=True`, nested
`tracknet`). ? `default_segmenter='yolo_hybrid'` HERE but config.json overrides to `gemini`. ?
`::IndexingConfig.segmenter_must_be_registered` `@field_validator` is a no-op pass-through ?
a typo'd segmenter validates fine, only fails later at `segmenter_registry.get(...)`).
97     - `::ValidationConfig` (`min_resolution_width=1280`, `min_resolution_height=720`, `min_fps=29.9`,
`min_blur_threshold=20.0`, `max_scene_cuts_per_minute=0.5`, nested `relevance`);
`::ValidationRelevanceConfig` (court HSV gates, `court_pixels_min_ratio=0.05`);
`::TrackNetConfig` (WSL invocation contract: `wsl_distro='Ubuntu'`, `repo_dir`,
`conda_env='TrackNetV4'`, `weights_path='', `queue_length=5`, `velocity_threshold=0.02`);
`::OutputConfig` (`default_highlights_name`, `db_name='sports_indexer.db'`,
`output_dir='output'` ? CLI `--output-dir` overrides it on the typed model in `main()`).

```

```

`::StitchingConfig`
(`begin_padding=0.5`,`end_padding=0.5`,`use_cache=False`,`min_shot_count=1`). `::Sport` Literal
of 5 sports. `::Config` = AppConfig` alias.
98     - `@backend/config/loader.py` ? load/save with defaults-merge.
99     - `::DEFAULT_CONFIG_PATH` = Path("config.json")` (? CWD-relative, NOT repo-root).
`::get_builtin_defaults` = `AppConfig().model_dump(...)` .
100    - `::load_config` ? precedence built-in-defaults -> file overrides. ? **shallow top-
level merge** (`{**defaults, **file_data}`): a section in config.json REPLACES the whole
defaults dict for that section (Pydantic then re-applies field defaults for missing leaves).
Missing/unreadable/parse-error file -> fully-defaulted AppConfig + `[CONFIG WARNING]` (non-
fatal).
101    - `::save_default_config` ? writes fresh `config.json`, **overwrites** if present
(used by setup.py/--setup). `::get_default_config` ? ? transitional plain-dict alias.
102    - `@config.json` ? on-disk config; mirrors `AppConfig`. ?
**`indexing.default_segmenter='gemini'`** diverges from model default `'yolo_hybrid'` ? file
wins at runtime.
103    - `@setup.py` ? `::init_default_config` ? delegates to
`backend.config.save_default_config` (single-source defaults; skips if config.json exists).
`::run_pip_install` GPU-aware (`nvidia-smi` -> cu124 else cpu). `::run_diagnostics` (Python?3.8,
ffmpeg/ffprobe on PATH, optional torch+CUDA).
104
105    ### 2.1 Segmenters (under pipeline/segmenters) ?
106    - `@backend/pipeline/segmenters/base.py`
107    - `::BasePlaySegmenter` ABC ? `__init__(db, config)`; abstract
`name`,`description`,`process_video`. $ `process_video(video_id, video_path, player_pool=None,
max_frames=None)` -> List[dict{id,start_time,end_time,duration,metadata}].
108    - `::SegmenterRegistry`,`::segmenter_registry` (singleton). ? **register a new
segmenter**: subclass `BasePlaySegmenter`, call `segmenter_registry.register(MyCls)` at module
bottom, AND import the module in `@backend/pipeline/segmenters/__init__.py` (registration is an
IMPORT side-effect). ? `register()` builds `cls(None, {})` just to read `.name` ->
name/description/`__init__` MUST NOT touch `self.db`/`self.config`. ? `list_available()` re-
instantiates + swallows exceptions into "No description available." (a throwing `description` is
masked).
109    - `@backend/pipeline/segmenters/__init__.py` ? imports all 5
(`motion,gemini,yolo_hybrid,tracknet_hybrid,wasb_hybrid`) = the de-facto registration manifest;
re-exports them + `segmenter_registry` via `__all__`.
110    - `@backend/pipeline/segmenters/wasb_hybrid.py::WasbHybridSegmenter` (name
`'wasb_hybrid'`) ? WSL WASB substrate (MIT; ships pretrained badminton weights, the LIVE
default). Imports `WasbRunner`/`WasbConfig` from `pipeline/detectors/wasb_runner`. Config
`indexing.wasb.velocity_threshold(0.02)` + shared
`dead_time_merge_gap/min_action_window_duration/ai_handoff_padding/ai_max_workers`. ? whole-
video single pass; reuses YOLO-named keys `log_yolo_candidates/store_yolo_candidates`; needs
`{PROVIDER}_API_KEY` (missing -> []); `detection_params` has NO `queue_length` (unlike
tracknet). ? divergent shot_count: `int(shot_count)` inside try with no `is None` short-circuit
(copy-paste drift).
111    - `@backend/pipeline/segmenters/tracknet_hybrid.py::TrackNetHybridSegmenter` (name
`'tracknet_hybrid'`) ? copy-paste twin of wasb_hybrid; imports `TrackNetRunner`/`TrackNetConfig`
from `pipeline/detectors/tracknet_runner`; `indexing.tracknet.velocity_threshold(0.02)`;
`detection_params` adds `queue_length`. ? any P3/P4 fix must be applied to BOTH (and arguably
yolo_hybrid).
112    - `@backend/pipeline/segmenters/yolo_hybrid.py::HybridYoloSegmenter` (name
`'yolo_hybrid'`, no YOLO remains ? torchvision detector + supervision ByteTrack; ADR-005).
`::PERSON_CLASS_ID=1` ? (ultralytics used 0; wrong value -> zero detections).
`::lazy_load_model` (lazy `torchvision` import so module/tests load without it),
`::extract_action_windows_from_chunk`,`::make_tracker` (per-chunk ByteTrack ? track IDs
chunk-local), `::DETECTOR_FACTORIES`/`::DEFAULT_DETECTOR='fasterrcnn_resnet50_fpn'`. ?
velocity at `effective_fps = fps/frame_skip` -> `yolo_velocity_threshold` coupled to
`frame_skip`; pipelined-vs-sequential P2 (prefetch ffmpeg overlaps sequential GPU);
`yolo_tracker`/`bytetrack.yaml` is dead/back-compat. ? `gemini_*` config keys are deprecated
aliases still honored.
113    - `@backend/pipeline/segmenters/gemini.py::GeminiPlaySegmenter` (name `'gemini'`) ?
pure-AI, no P2. ? chunks only if `duration>chunk_duration`; `max_workers=5` hardcoded;
`max_segment_duration=90.0` hardcoded; `debug_duration` silently truncates; `parse_time` accepts
colon timestamps (warns); does NOT persist candidate segments (only hybrids feed the fine-tuning
table).

```

```

114     - `@backend/pipeline/segmenters/motion.py::MotionPlaySegmenter` (name `motion`) ?
pixel absdiff baseline. ? **metadata + serve/smash/foul/winner events are RANDOM**
(random.sample/randint/uniform/choice); only segmenter using `db.add_event` and passing
server/receiver; uses `print()`; honors `max_frames`.
115
116     ### 2.2 Detectors & windowing (WSL-quarantined; ABC abstraction complete)
117     - `@backend/pipeline/detectors/base.py` ? DetectorRunner ABC (healthcheck() ->
Tuple[bool,str] never-raises; run_predict(...) -> Optional[str] CSV path). Contract deliberately
matches the pre-existing tuple returns and call sites in the hybrids. Rich module docstring
includes "how to add a new runner", Linux-native path notes (the de-risk goal), and guidance on
threading health into reports. Re-exported from `detectors/__init__.py`. This finishes the WSL
de-risk seam (low-regret 6 / review item 1).
118     - `@backend/pipeline/detectors/tracknet_runner.py`
119     - `::TrackNetConfig` (+`.from_indexing_cfg`; ? key drift `wsl_distro` JSON -> `distro`
field), `::TrackNetRunner` ? runs TrackNet `predict.py` in WSL.
120     - `::to_wsl_mnt_path` (`C:\x`->`/mnt/c/x`; ? POSIX/~` passthrough, UNC unhandled),
`::TrajectoryPoint`(frame,visible,x,y).
121     - `::parse_trajectory_csv` @staticmethod** + `::trajectory_to_action_windows`
@staticmethod** = MODEL-AGNOSTIC brain, re-exported verbatim by WasbRunner. ? `weights_path`
defaults `` -> `run_predict` returns None; predict.py only `.mp4/.avi` + hard-names
`<stem>_predict.csv`; `visible==False` resets prev (occlusion gap breaks track); `fps<=0->30`,
`frame_width<=0->1280` silent; `track_id_count` hardcoded 1; `timeout_sec=0` -> NO timeout
(hangs forever). `::stage_video` "OK" sentinel.
122     - `@backend/pipeline/detectors/wasb_runner.py::WasbRunner` (inherits DetectorRunner)
(+`::WasbConfig`, from `indexing.wasb`; ? `weights_path` HAS a real default
`wasb_badminton_best.pth.tar` ? no required guard) ? stages `wasb_infer.py` (`::INFER_WIN`) +
video into WSL, runs in `wasb` env. `::parse_trajectory_csv`/`::trajectory_to_action_windows` =
`staticmethod(TrackNetRunner.*)` (explicit rebind = the plugin-equivalence point). Output hard-
named `<stem>_wasb.csv`. ? `wasb_infer.py` re-staged each run (editing Windows copy inert until
staged); `wsl_bash` duplicated (drift risk). Both runners now satisfy the DetectorRunner ABC
for future swappability.
123     - `@backend/pipeline/detectors/wasb_infer.py` (WSL) ? `::run` core. ? imports WASB-SBDT
repo modules + torch ? NOT importable on Windows. Reboot-resumable detector cache:
`::DET_FILE='det_raw.jsonl'`, `::MANIFEST_FILE='manifest.jsonl'`, `::CACHE_VERSION=1`,
`::load_and_clean_cache` (reads longest CONTIGUOUS prefix where `w==line index`, rewrites to
heal crash-truncated tail), `::manifest_compatible` (? `frames_dir` abspath'd but `weights`
compared raw), `::atomic_write_text`/`atomic_write_trajectory`, `::dets_to_tracker` (xy MUST be
np.array), `::build_cfg` (`runner.gpus=[0]` hardcoded), `::extract_frames` (cv2-PNG SLOW),
`::default_cache_dir` (`<stem>_wasbcache` next to out). ? GPU pass resumable; CPU tracker
STATEFUL -> always fully re-run; `--fresh` ignores cache.
124     - `@backend/pipeline/detectors/rally_gate.py` ? Gate A net-crossing post-filter (pure,
no I/O/model). `::estimate_play_axis` (larger-spread axis; ? ties -> 'x'),
`::count_net_crossings` (sign changes vs median net; deadband jitter zone),
`::gate_windows`/`::apply_gate` (`min_crossings=2` default; kills single-toss service feeds).
`::GatedWindow`(start,end,crossings,visible_points,kept). ? window contract here is
`(start,end)` tuples NOT the rich window dict ? caller must adapt. Consumed by calibrate_local /
distill_local / export_wasb_segments.
125
126     ### 2.3 Eval & calibration (all pure core; I/O in CLI drivers)
127     - `@backend/eval/metrics.py` ? scoring primitive (THE spine). `::interval_iou`,
`::match_intervals` (greedy, deterministic), `::score`->`Score`(`.as_dict()` = wire shape),
`::pr_curve`, `::Match`, `::MatchResult`. ? greedy not Hungarian; empty matches -> 0.0 (not NaN)
for mean_iou/boundary errors.
128     - `@backend/eval/annotations.py` ? generic GT loader. `::load_annotations` (`start,end`
CSV; RAISES on bad rows/`start>=end`), `::parse_timestamp` (decimal or MM:SS/HH:MM:SS),
`::write_scaffold`.
129     - `@backend/eval/calibration.py` ? overfit guard. `::cross_validate` (within-video
k-fold; ? defaults `metric='recall` deliberately), `::leave_one_video_out` (honest cross-video;
needs ?2 labelled videos), `::improvement_report` (OPTIMISTIC ? selected on full GT),
`::select_best`, `::kfold_indices`, `::pct_improvement` (? from-zero -> `inf`), `::evaluate`. $
`PredictFn = Callable[[Config], List[Interval]]` = the plug-in seam. ? `generalization_gap
>~0.1` = overfit alarm.
130     - `@backend/eval/calibrate_wasb.py` ? WASB tuning CLI. `::BASELINE=(0.02,2.0,2.0)`,
`::default_grid` (45 cfgs), `::make_predict_fn` (parses trajectory once, closes over
`trajectory_to_action_windows`), `::load_golden_gts` ($`start_time,end_time` cols; ? SILENTLY

```

```

skips bad rows ? opposite of annotations.load_annotations). ? `{load_golden_gts,
make_predict_fn, BASELINE}` imported by 4 downstream drivers (calibrate_local,
export_wasb_segments, gemini_refine, audio/e1_probe).
131 - `@backend/eval/calibrate_local.py` ? distill Gemini->local thresholds (windowing +
net-crossing gate, no cloud at runtime). `::local_grid` (180 cfts), `::make_local_predict_fn`
(adds `rally_gate.apply_gate` when `min_crossings>0`), `::build_videos` (? `SystemExit` if a
labels file yields no rallies), `::run`/`::main`. $
`LocalConfig=(velocity,merge,min_dur,min_crossings)`, `BASELINE_LOCAL=(0.02,2.0,2.0,0)`;
manifest JSON shared with distill_local.
132 - `@backend/eval/export_wasb_segments.py` ? tuned cfg -> golden/slicer CSV + comparison
+ optional clip cut. `::TUNED=(0.05,1.0,1.0)` (? from ONE video),
`::SEG_FIELDS`, `::COMP_FIELDS`, `::build_comparison`, `::seconds_to_mmss`, `::maybe_slice`
(lazy-imports rally_slicer). `--slice` requires `--video`. ? `write_segments_csv` output loads
straight into `rally_slicer.load_rally_labels` AND re-imported by gemini_refine ? schema is
load-bearing.
133 - `@backend/eval/batch.py` ? corpus aggregation (pure; orchestration is in
run_phase3_eval). `::aggregate` (micro pool TP/FP/FN, TP-weighted mean_iou, macro F1),
`::default_stages_for` (`auto`; `::FINAL_ONLY_SEGMENTERS={motion,gemini}`),
`::resolve_video_path` (normalizes collector `./data\..` windows paths), `::format_aggregate`.
134 - `@backend/eval/harness.py` ? DB-pred scorer (no re-run).
`::STAGES=('candidates','final')`, `::get_predictions`
(candidates->`query_candidate_segments(detector=)`; final->`query_play_segments({})`); ? unknown
stage raises ValueError; reused by regression.py),
`::evaluate`->`::EvalReport`(`::StageReport`)->`::report_to_dict` (input contract for
batch.aggregate), `::format_report_text`.
135 - `@backend/eval/training.py` ? learned-segmenter data layer. `::YOLO_FEATURE_NAMES`
(5-d), `::extract_yolo_features` (? reads `row['stats']` dict from DB; missing fields default
0.0), `::label_candidates` (same IoU `match_intervals` as evaluator),
`::build_training_set`->(X,y,intervals); `feature_fn` is the substrate plug-in seam.
136 - `@backend/eval/classifier.py::LogisticRegression` ? numpy-only logreg (no sklearn).
`fit/predict_proba/predict/to_dict/from_dict/save/load`. ? `class_weight='balanced'` default;
stores feature standardization (mean/std) ? inference MUST use same feature ORDER; `_sigmoid`
clips z to [-30,30]; `lr=0.1,epochs=1000,l2=1e-3`; JSON "model":"logreg" tag.
137 - `@backend/eval/eval_partial_labels.py` ? score pipeline vs **PARTIAL human labels**
(PR #60). `::restrict_to_window` (drop preds outside `[0, last GT end]` so un-labeled-tail
detections aren't FPs ? THE partial-label methodology), `::run` (baseline/+gate/tuned/+gate ops
table + per-rally diff via `export_wasb_segments.build_comparison`), `::sweep`/`--sweep` (grid
vs labels; ? fit-on-self = optimistic), `::DEFAULT_GRID`, `::TUNED=(0.05,1.0,1.0)`. ?
`window_end` inferred from `max(GT end)` is WRONG if a video is labelled PAST its last rally ?
P1 fix = collector persists `labeled_window` in the manifest. Reuses calibrate_wasb + rally_gate
+ metrics.
138 - `@backend/eval/rally_seq_proto.py` ? **PROTOTYPE** learned per-FRAME rally segmenter =
owned-model **Gen-0 seed** (PR #61). 6 features over the WASB trajectory (`vis_rate, spd_mean,
spd_max, cross_rate, x_std, y_std`) -> `classifier.LogisticRegression`, honest LOVO.
`::build_frame_arrays`, `::features` (scipy `uniform/maximum_filter1d` windowed; ? speed NOT
fps-normalized yet ? 30 vs 60fps), `::labels_for`, `::probs_to_segments`
(smooth->threshold->merge/min-dur), `::roc_auc` (rank-based, no sklearn), `::run` (per-video AUC
+ LOVO-threshold F1 + oracle-threshold F1). ? DIRECTIONAL only (n?5, linear, threshold-transfer
unsolved); held-out per-frame AUC 0.83?0.92, **beats the heuristic on multi-court footage**
(Kushagra/gBaaddy). Manifest `output/human_lovo_manifest.json` (the n=5 human corpus) shared w/
calibrate_local.
139 - ? **MULTI-COURT / background-game confound (2026-06-08 finding):** WASB tracks EVERY
shuttle in frame incl. other courts; videos with background games (Kushagra, gBaaddy) have high
dead-time shuttle visibility (41?57% vs 18?22% clean) ? the velocity-windowing heuristic can't
find clean rally boundaries (F1 0.16/0.28 vs ~0.75 on single-court). Diagnostic = in-rally+dead-
time visibility "contrast" (?3.6x heuristic works, ?1.9x fails). Lever = **foreground-court
isolation** (spatial gate to the prominent court, needs real court-region detection). Academy =
multi-court = the target env. See `docs/archives/quality-iterations/2026-06-multivideo-lovo/`.
140 - `@backend/eval/distill_local.py` ? distill Gemini -> a LEARNED classifier (beyond
calibrate_local's threshold ceiling). `::FEATURE_NAMES` (5 fps/res-INVARIANT:
duration,peak,mean,active_ratio,net_crossings), `::build_candidates` (recall-first proposal),
`::predict_rallies` (proba filter then `gemini_refine.merge_overlaps(gap_tol=1.0)`),
`::run`/`::main`. consts `PROPOSAL=(0.005,1.0,0.5)`, `BASELINE_LOCAL=(0.02,2.0,2.0)`. ? leave-
one-video-out needs ?2 videos; final model trained on all.
141 - `@backend/eval/gemini_refine.py` ? refine WASB candidates with Gemini over SHORT

```

```

padded clips (precision lever; eval/tuning, NOT the live pipeline).
`::merge_overlaps(gap_tol=0.0)` (reused by distill_local with gap_tol=1.0), `::refine_window` (?
FRESH genai client per worker ? shared throws socket errors), `::full_video_rallies`,
`::run`/`::main` (`--full-video` skips `--trajectory`). ? requires `GEMINI_API_KEY` (raises
SystemExit); imports `BASELINE` from calibrate_wasb +
`TUNED`/`write_segments_csv`/`seconds_to_mmss` from export_wasb_segments.
142 - `@backend/eval/regression.py` ? CI gate. `::DEFAULT_TOLERANCE=0.05`,
`::DEFAULT_BASELINE_PATH=output/regression_baseline.json`, `::regress`,
`::regress_with_provider` (GT via ? `annotation_registry.get(tier)`;
"file"->FileProvider/"vendor"->HumanVendor/"auto"->oracle),
`::save_baseline`/`::load_baselines`, `::RegressionResult`. ? imports `annotation_registry` from
**`backend.annotations`** NOT `backend.eval.annotations`; gate fires on precision OR recall drop
(F1 reported, not gated); no baseline -> `regressed=False` (silent pass).
143
144     ### 2.3a Audio (E1 probe; NOT adopted for fusion ? ADR-007)
145 - `@backend/pipeline/audio/hit_detector.py` ? classical-onset shuttle-hit detector. Pure
numpy + stdlib `wave` (NO scipy); ffmpeg only for extraction. `::HitEvent`(t,strength),
`::extract_audio` (ffmpeg -> mono 16k PCM WAV), `::load_wav`, `::onset_envelope` (pure-numpy
STFT spectral flux; `band=(lo,hi)` band-limit), `::detect_hits` (adaptive `mean+k.std`;
`k`=sensitivity knob), `::detect_hits_in_video`. ? separate signal ? WASB never modified by
this; per audio-e1-findings recall 100% but discrimination ~2.2x, band-pass null -> NOT adopted
(PR #35).
146 - `@backend/eval/audio/e1_probe.py` ? E1 GO/NO-GO diagnostic (no training/fusion).
`::cluster_hits` (groups thwacks; ? `min_hits>=2` so a lone service-feed hit isn't a rally),
`::run` (discrimination ratio + recall recovery + audio-only P/R/F1), `::main`. reuses
`hit_detector` + `calibrate_wasb.{load_golden_gts,make_predict_fn,BASELINE}` + `metrics`.
147
148     ### 2.4 Stitcher / tools / utils
149 - `@backend/pipeline/stitcher/compiler.py::VideoCompiler` ? `compile_highlights(raw,
segments:[(start,end)], out_name)`. Per-clip re-encode -> `-f concat -c copy`. `::parse_time`
(mirrors genai.parse_time; ? unparseable -> 0.0), `::get_video_duration`. ? if duration probe
fails (0.0) NO boundary validation; `begin_padding`(0.5)/`end_padding`(1.0) defaults; temp clips
in a TemporaryDirectory under output_dir; raises ValueError/FileNotFoundError/RuntimeError; uses
`print()`. ? two-stage codec contract: clips re-encoded with IDENTICAL settings precisely so
final concat can stream-copy ? change a preset and concat may break.
150 - `@backend/tools/rally_slicer.py` ? standalone CLI. `::compute_clip_windows` (pure,
unit-tested), `::load_rally_labels` (golden schema; ? silently skips degenerate rows),
`::slice_rallies`, `::ClipWindow`(rally:str), `::BEGIN_PADDING=0.5`/`::END_PADDING=1.0` (?
duplicate compiler's by copy). `read_time` -> backend/eval/annotations.parse_timestamp.
151 - `@backend/utils/video.py` ? `::get_video_duration` (ffprobe; ? 0.0 on failure =
"unknown"), `::extract_video_chunk(...,reencode=False,scale_width=None)`. ? copy mode cuts at
nearest keyframe (drift) ? pass `reencode=True` for short/accurate clips; `scale_width` requires
`reencode=True`.
152 - `@backend/utils/geometry.py` ? `::get_velocity_normalised` (fraction of frame-width/s;
? raw-px fallback if width<=0; used by yolo_hybrid), `::calculate_iou` (? +1 px convention
inflates IoU), `::get_center/get_distance/get_velocity`.
153
154     ### 2.5 Validators ?
155 - `@backend/pipeline/validators/base.py::VideoValidator` ABC, `::ValidationResult`
(pydantic: validator_name,passed,message,details), `::ValidationRegistry` (ordered LIST ?
registration order = exec order; ? no dedup), `::registry` (singleton). $ `validate(video_path,
config, context)` -> `ValidationResult`. `::run_all_with_context` -> (results, context) (wraps
each validate in try/except -> failed result not abort); `::run_all` discards context. ? **add a
validator**: subclass + `registry.register(MyValidator())`.
156 - `@backend/pipeline/validators/__init__.py` ? registers in EXEC order: FormatValidator,
ResolutionFPSValidator, PTSContinuityValidator, SceneCutDensityValidator, BlurValidator,
RelevanceValidator (? producers before consumers ? format_check first; ? import mutates global
registry as side effect).
157 - `@backend/pipeline/validators/format.py::FormatValidator` (`format_check`) ?
**PRODUCER** of `context['ffprobe_meta']`; ? MUST run before resolution_fps; container check is
substring `mp4` (so MOV-family passes); external `ffprobe`.
158 - `@backend/pipeline/validators/resolution_fps.py::ResolutionFPSValidator`
(`resolution_fps_check`) ? CONSUMES `ffprobe_meta` (? hard-fails "Run format_check first" if
absent). `min_fps` 29.9; min res 1280x720.
159 - `@backend/pipeline/validators/pts_continuity.py::PTSContinuityValidator`

```



```

(pts_continuity_check`) ? re-probes (no context). ? 10.0s keyframe-gap threshold hardcoded
(comment says 8.0s ? drift, NOT config-driven); any backward jump fails; `<2` keyframes PASSES
(too short).
160 - `@backend/pipeline/validators/scene_cut.py::SceneCutDensityValidator`
(`scene_cut_check`) ? OpenCV. ? `cut_threshold=0.6` hardcoded (NOT config); samples ~2fps, caps
at 100 SAMPLED frames (~first 50s); `max_scene_cuts_per_minute` (0.5) is the config fail
threshold.
161 - `@backend/pipeline/validators/blur.py::BlurValidator` (`blur_check`) ? Laplacian var <
`min_blur_threshold` (default 100.0 in-file, 20.0 in models/config) = fail; samples 15 frames. ?
magic absolute, not resolution-normalized.
162 - `@backend/pipeline/validators/relevance.py::RelevanceValidator` (`relevance_check`) ?
HSV court-green ratio (? purely color, no geometry despite name); lazy `import backend.sports`;
3-tier config precedence (`validation.relevance` -> sport profile -> hardcoded); ? unknown sport
-> empty cfg, falls back to defaults, does NOT fail.
163
164 ### 2.6 Sports / Providers / Annotations ?
165 - `@backend/sports/base.py::SportProfile` ABC
(`name`, `get_prompt`, `get_relevance_config`);
`@backend/sports/badminton.py::BadmintonSportProfile` (the Gemini prompt + HSV relevance cfg);
`@backend/sports/__init__.py::sport_registry` (+`::SportRegistry`, auto-registers badminton; ?
`register` instantiates profile to read `name`). ? **add a sport**: subclass `SportProfile`,
`sport_registry.register(MyProfile)`. $ prompt emits JSON array `{start_time(float DECIMAL SEC),
end_time, shuttle_exchange_count(int), shot_count_confidence('Low'|'Medium'|'High')}`. ? prompt
actively defends against MM:SS; `get()` returns a fresh instance each call.
166 - `@backend/providers/base.py::AIVideoProvider` ABC (`__init__(api_key, model_name)`);
`@backend/providers/gemini.py::GeminiProvider` (lazy `::client` property; `::MAX_UPLOAD_MB=500`;
upload->poll->generate JSON temp 0->delete); `@backend/providers/__init__.py::provider_registry`
(+`::ProviderRegistry`, auto-registers `gemini`). ? **add a provider**: subclass,
`provider_registry.register('name', MyCls)` (key passed EXPLICITLY); `get(name, api_key,
model_name)`. $ `analyze_video(path, prompt, chunk_duration=0.0, label='')` -> (segments:list,
metrics:dict{upload_time, process_time, gen_time}). ? blocking 10s poll; returns `([], metrics)`
on ANY failure (empty != error); oversize clip silently skipped; orphaned remote file possible
if process dies mid-run.
167 - `@backend/annotations/base.py::AnnotationProvider` ABC (`submit/poll/fetch/annotate`,
`::Interval` (= metrics.Interval), `::QUEUED/RUNNING/DONE/FAILED`/`::STATUSES`,
`::AnnotationProviderRegistry` (decorator-friendly `register` returns the class; ? keys off
`provider_cls().name`, falls back to `__name__` if no-arg ctor throws; `list_available()` ->
LIST not dict), `::annotation_registry`. ? **add a tier**: subclass (define `name`),
`annotation_registry.register(MyCls)` at module bottom + import in `__init__.py`. $ `fetch` ->
List[Interval]` sorted; contract says RAISE (not `[]`) on unavailable.
168 - `@backend/annotations/file_provider.py::FileProvider` (`file`; job_id IS the
resolved CSV path; `_resolve` -> `/.csv`; ->
`backend.eval.annotations.load_annotations`; ? `guideline`/`sport` ignored).
169 - `@backend/annotations/automated_provider.py::AutomatedProvider` (`auto`; GS-4 stub;
injectable `::Labeler`, in-process `_jobs` store, `::warn_if_same_lineage` ORACLE RULE ? only
warns, does not block; default `_not_configured` RAISES NotImplementedError). ? get via
`annotation_registry.get('auto', labeler=fn, lineage='...')`; registered as `auto` only
because no-arg ctor succeeds.
170 - `@backend/annotations/__init__.py` ? ? `noqa F401` imports of
file_provider/automated_provider are load-bearing (removing de-registers providers).
171
172 ### 2.7 Infrastructure (DB) ?
173 - `@backend/infrastructure/database.py::Database` ? SQLite, 4 tables
(`videos/play_segments/events/candidate_segments`), `PRAGMA user_version` migration ladder in
`::init_db` (currently =1, `version<1` creates `candidate_segments`). `::get_connection` (?
re-issues `PRAGMA foreign_keys=ON` per connection ? else CASCADE/SET NULL silently skipped).
Methods: `add_video`, `add_play_segment`, `add_event`, `add_candidate_segment`, `link_candidate_to_s
egment`, `query_candidate_segments(detector=, only_unlabeled=)`, `query_play_segments`, `get_video`,
`clear_video_data`. ? candidate_segments is the detector-agnostic fine-tuning extension point
(common fields columns, detector-specifics in `stats`/`detection_params` JSON; index
`idx_candidates_video_detector`). ? `add_video` is `INSERT OR REPLACE` -> cascade-DELETES
dependent segments on re-ingest; events column is `type` not `event_type`; `query_play_segments`
filters use truthiness so `duration_min=0`/empty-string = "no filter"; write helpers swallow
exceptions.
174

```

```

175     ### 2.8 Reporting / observability ? (per-video, soft-dep `obsreport`)
176     - `@backend/reporting/__init__.py` ? `::emit_indexer_report(*, db, video_id, video_path,
config, val_results, val_context=None, segmenter_requested=, segmenter_actual=,
was_reingest=False, segmentation_ran=True)` ? single entrypoint, keyword-only. Normalizes typed
config (`config.model_dump(...)` if `hasattr model_dump`); short-circuits to None if
unavailable/disabled; pulls play_segments/candidates from `db` (read-only
`query_play_segments(vid, {})`/`query_candidate_segments(vid)`); delegates to
`harvest.record_run`; writes via `rec.write(output_dir=config["report"]["output_dir"])`. ? wraps
everything in bare `except Exception: return None` ? **NEVER raises** (designed for `main.py`'s
`finally:`); writes `

```

```

195     $ **play_segment dict** (`db.add_play_segment` / `query_play_segments`):
196     `{id:int, video_id, start_time:float, end_time:float, duration:float(end-start), server,
receiver, metadata:dict, events:[...]}`. Hybrid P4 metadata `{shot_count,
shot_count_confidence:'Unknown', detector:f'{family}-hybrid-{model}'}`; gemini detector = bare
`model_name`; motion adds random `intensity/avg_speed_kmh`.
197
198     $ **AI segment dict** (provider output, consumed by hybrid P3 / gemini): `{start_time,
end_time, shuttle_exchange_count?, shot_count_confidence?}` (offset by chunk/window start;
`_candidate_id` stamped internally).
199
200     $ **GT / golden CSVs** (TWO conventions ? do not conflate):
201     - generic (`annotations.load_annotations`): `start,end` 2-col; optional header; decimal
or MM:SS/HH:MM:SS; RAISES on bad row.
202     - golden/slicer (`calibrate_wasb.load_golden_gts`, `rally_slicer.load_rally_labels`,
`export.SEG_FIELDS`): `rally_number,start_time,end_time[,duration,start_mmss,end_mmss]`;
SILENTLY skips bad rows.
203
204     $ **det_raw.jsonl** (one window/line): `{ "w":<window_index>,
"f":[[frame_id,[x,y,score,scale],...]], ...]}`. Lines MUST be contiguous (line N has `w==N`);
first violation/JSONDecodeError -> truncate-heal.
205     $ **manifest.json** (cache): `{version, frames_dir(abspath), weights, sport, frames_in,
frames_total, windows_total, windows_done}`. Reuse only if ALL match + `version==CACHE_VERSION`.
206     $ **collector manifest.json** (eval): `{eval_sets:[{video_id(label), local_path, csv,
...}]}`; ? DB key = file MD5, manifest `video_id` is a label.
207
208     $ **eval primitives**: `Interval=Tuple[float,float]` (start<end).
`Match{pred_index,gt_index,iou}`. `MatchResult{matches, unmatched_pred_indices(=FP),
unmatched_gt_indices(=FN)}`. `Score{iou_threshold,num_pred,num_gt,true_positives,false_positives
,false_negatives,precision,recall,f1,mean_iou,mean_start_error,mean_end_error}.as_dict()`.
209
210     $ **WSL detector candidate** (in-WSL, model<->tracker): `{ 'xy':np.array([x,y]),
'score':float, 'scale':int}`; plain JSON form `[x,y,score,scale]`. ? `_dets_to_tracker` MUST
rebuild `xy` as np.array (tracker does `np.linalg.norm`).
211
212     $ **HitEvent** (audio, `pipeline/audio/hit_detector`): `{t:float(sec), strength:float}`
? onset-envelope peak.
213
214     $ **observability report** (`reporting/harvest.record_run` -> `RunRecorder`): stages
`validation/segmentation/ai_handoff` + highlights + verdict ? {pass,pass_with_warnings,fail};
flags `{code,faq_ref}`; `video_meta.fps_source` ? {probed,fallback}. ? stage/metric/detail names
are coupled to `spec.yaml::rules[].when[].path` ? pairs that MUST stay in sync:
`ai_handoff.details.{api_key_present,ai_based}`, `segmentation.metrics.segment_count`, `segmenta
tion.details.{segmenter_actual,matches_config_default,segmenter_explicitly_requested,was_reinges
t}`, `validation.status`, `video_meta.fps_source`.
215
216     ---
217
218     ## 4. WASB WSL CONTRACT (external dep `~/models/WASB-SBDT`, MIT; NOT in this repo ?
claims below are unverified against external source, only the repo-side caller wasb_infer.py was
checked)
219     - Env: WSL conda `wasb`, cwd `~/models/WASB-SBDT/src`. Hydra
`compose(config_name='eval', overrides=[dataset=<sport>, model=wasb,
detector.model_path=<weights>, runner.device=cuda, runner.gpus=[0]])`. `gpus=[0]` because box is
single-GPU. (overrides VERIFIED in @backend/pipeline/detectors/wasb_infer.py::build_cfg.)
220     - `model=wasb` -> `name=hrnet`, `frames_in=3`, `frames_out=3`, input 512x288 (WxH),
ImageNet normalize; affine resize maps heatmap coords -> original-frame coords. (model dims read
from cfg["model"] at runtime in wasb_infer.run.)
221     - `detectors/detector.py::TracknetV2Detector.run_tensor(imgs, affine_mats)` -> (results,
hms_vis). `results[bid][eid]` = list of `{xy:np.array (ORIGINAL coords), score, scale}`.
(caller uses `detector.run_tensor(imgs, trans)` returning `(batch_results, _)` ? consistent.)
222     - `trackers/online.py::OnlineTracker` ? STATEFUL, sequential per `update()` (assumes
every frame fed once, in order, no gaps). `update(frame_dets)->{x,y,visi,score}`; `refresh()`
resets. `det['xy']` MUST be np.array. Full ordered re-run is deterministic. (caller calls
`build_tracker(cfg)`, `tracker.refresh()`, `tracker.update(...)` returning dict with `visi/x/y`
? consistent.)

```

```

223     - `dataloaders/dataset_loader.py::ImageDataset(cfg, samples, input_wh, output_wh,
transform, seq_transform, is_train)`; `samples=[{frames:[paths], annos:[{frame_path,
center:Center(is_visible,x,y)}]]`. (caller constructs exactly this ? VERIFIED in
wasb_infer.run.)
224     - Weights `pretrained_weights/wasb_badminton_best.pth.tar` ? NOT committed, academic-
data-trained -> ? confirm terms before COMMERCIAL deploy. monotrack weights = noncommercial,
avoided.
225
226     ---
227
228     ## 5. CROSS-REPO (sibling private repo `sports-data-collector` ? external, unverified
here)
229     - Indexer keys video by **FILE MD5**; collector keys by human `video_id`. ? re-
encoding/re-downloading changes bytes -> MD5; acquisition MUST precede processing; manifest
tracks the exact file used.
230     - `python -m src.cli.curate export` (collector) -> `//.csv`
(indexer `start,end` format) + `manifest.json` -> consumed by `@backend/eval/batch.py` +
`@run_phase3_eval.py`.
231     - Collector also emits golden rally labels (`rally_number,start_time,end_time,...`) ->
consumed by `@backend/eval/calibrate_wasb.py` / `export_wasb_segments.py` /
`@backend/tools/rally_slicer.py`.
232
233     ---
234
235     ## 6. GOTCHAS (cross-cutting footguns)
236     - **MD5 keying:** `video_id` = MD5 of WHOLE file via the single
`@backend/utils/hashing.py::compute_video_id` helper (64KB chunks), imported by main + all
run_*.py. Manifest `video_id` is a label, NOT the DB key.
237     - **Config access:** `backend/main.py` and all four `run_*.py` load via
`backend.config.load_config` (typed `AppConfig`). The legacy `.get()` dict-compat shim remains
for validators/segmenters. Remaining literal: a defensive `getattr(..., 'motion')` fallback in
`main.py` that only fires if `global_config`/.indexing` is `None` (the ASGI-import edge case).
238     - **Segmenter default:** single source of truth = model
`IndexingConfig.default_segmenter='yolo_hybrid'`; `config.json` may override at runtime (ships
`gemini`). run_*.py read the model value (no literals); `main.py` keeps the defensive `motion`
getattr fallback noted above.
239     - **`AppConfig.get` must return dict sections, not models:** the legacy
`.get("indexing",{}).get(...)` chains everywhere depend on it; the `extra="allow"` model also
silently keeps typo'd config keys (no rejection).
240     - **`output_dir`:** a real field on `OutputConfig` (default `"output"`); the CLI
`--output-dir` overrides it on the typed model in `main()`, and `compile_highlights` reads the
same `global_config.output.output_dir` ? honored end-to-end.
241     - **Gemini free-tier / API-key dead path:** hybrids + gemini return `[]` (not error) if
`{PROVIDER}_API_KEY` missing OR WSL healthcheck fails ? a silent no-op. `analyze_video` also
returns `[]` on real failure -> empty != "no rallies". The reporting
`silent_provider_noop`/`ai_empty_vs_no_rallies` rules exist precisely to disambiguate this in
the report.
242     - **Reporting never breaks a run:** `emit_indexer_report` swallows all exceptions
(designed for `main.py` `finally:`) and is emitted even on validation failure. `obsreport` is a
SOFT dep ? absent -> no-op.
243     - **spec.yaml <-> harvest.py coupling:** rule `path` strings dotted-navigate harvest's
stage/metric/detail names; rename a detail in harvest.py and the matching rule silently stops
firing. Add rules in spec.yaml, not code.
244     - **fps / frame_width must be PROBED not assumed:** `_probe_fps_width` / runner
fallbacks (30.0, 1280.0) silently mis-scale velocity thresholds; calibrate_wasb CLI bakes
`59.94/1920`. The report's `fps_fallback_used` flag surfaces this.
245     - **WSL `/mnt` can't see Google Drive / network mounts:** stage video into native WSL fs
(`~/clips`) ? `/mnt/c` reads are slow and Drive-backed paths invisible.
246     - **ffmpeg NOT in WSL** by design: all ffmpeg/ffprobe run Windows-side (must be on
PATH); WSL only runs the GPU model. `extract_frames` uses cv2 PNG (slow) ? prefer host ffmpeg +
`--frames_dir`. ? Don't import `wasb_infer` from Windows code (it imports torch + WASB repo
modules).
247     - **GPU float nondeterminism:** detector pass is checkpointed/resumable but not bit-
reproducible; the CPU tracker re-run over the ordered cache IS deterministic.
248     - **`output/` is gitignored:** DB, baselines, cache, reels live there. Predictions MUST

```

exist in `output/<db\_name>` before any eval/regression.

249 - ****Validator ordering coupling:**** exec order is Format -> ResolutionFPS -> PTSContinuity -> SceneCut -> Blur -> Relevance; `resolution\_fps\_check` hard-fails unless `format\_check` ran earlier (shared `context['ffprobe\_meta']`). PTS 10.0s gap + scene-cut 0.6 correlation are NOT config-overridable.

250 - ****`add\_video` INSERT OR REPLACE + FK CASCADE**** -> re-ingesting a `video\_id` cascade-deletes its segments. Use `clear\_video\_data` deliberately. `\_get\_connection` MUST re-issue `PRAGMA foreign\_keys=ON` per connection.

251 - ****Twin segmenters:**** `wasb\_hybrid.py` / `tracknet\_hybrid.py` are copy-paste; fixes to P3/P4 must land in both. ? subtle drift: wasb's shot_count `int(None)`-via-TypeError fallback vs tracknet's `is None` check.

252 - ****Two "annotations" modules:**** `backend.eval.annotations` (GT CSV loader) vs `backend.annotations` (provider registry). regression.py imports the latter.

253 - ****`calibrate\_wasb.{load\_golden\_gts, make\_predict\_fn, BASELINE}`** is a de-facto shared API** for calibrate_local, export_wasb_segments, gemini_refine, e1_probe ? changing its golden-CSV schema or `BASELINE` ripples widely. `gemini\_refine.merge\_overlaps` reused by `distill\_local`; `harness.get\_predictions` reused by `regression.regress`.

254 - ****Hardcoded tuned constants:**** `calibrate\_wasb.BASELINE`, `export.TUNED`, slicer/compiler padding ? duplicated by copy, not import; re-tuning won't propagate.

255 - ****`timeout\_sec=0` = no timeout:**** a hung WSL/GPU call blocks forever.

3. user (2026-06-10T05:59:37.508Z)

/c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/automated_provider.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/base.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/file_provider.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/__init__.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/config/loader.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/config/models.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/config/__init__.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/annotations.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/audio/e1_probe.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/audio/__init__.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/batch.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/calibrate_local.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/calibrate_wasb.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/calibration.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/classifier.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/distill_local.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/eval_partial_labels.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/export_wasb_segments.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/gemini_refine.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/harness.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/metrics.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/rally_seq_proto.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/regression.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/training.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/eval/__init__.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/infrastructure/database.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/main.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/audio/hit_detector.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/audio/__init__.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/base.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/rally_gate.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/tracknet_runner.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/wasb_infer.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/wasb_runner.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/__init__.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/base.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/gemini.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/motion.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/tracknet_hybrid.py
 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/wasb_hybrid.py

```

/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/yolo_hybrid.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/__init__.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/stitcher/compiler.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/base.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/blur.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/format.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/pts_continuity.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/relevance.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/resolution_fps.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/scene_cut.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/__init__.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/providers/base.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/providers/gemini.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/providers/__init__.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/reporting/harvest.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/reporting/__init__.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/results.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/sports/badminton.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/sports/base.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/sports/__init__.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/tools/export.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/tools/rally_slicer.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/tools/__init__.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/geometry.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/hashing.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/video.py
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/__init__.py
---
/c/Users/avidu/Projects/badminton-highlight-indexer/backend/reporting/spec.yaml

```

4. user (2026-06-10T05:59:44.386Z)

```

256     - **Stitch concat assumes identical encode:** per-clip re-encode then `-c copy` concat;
if probe fails (0.0) no boundary validation.
257     - **Audio NOT in the live pipeline:** `pipeline/audio/hit_detector` +
`eval/audio/e1_probe` are a diagnostic probe (ADR-007); recall 100% but discrimination ~2.2x ->
not adopted for fusion.
258
259     ---
260
261     ## 7. RAMP PATHS (to do X -> open Y::symbol)
262     | Task | Open |
263     |---|---|
264     | Change a config default / add a config field | `@backend/config/models.py` (single
source of truth) -> regenerate via `setup.py`/`save_default_config`; ? config.json may override
(default_segmenter) |
265     | Read config in new code | `from backend.config import load_config, AppConfig`
(`@backend/config/__init__.py`); attribute access, NOT raw `json.load` |
266     | Tune candidate windowing (velocity/merge/min-dur) |
`@backend/pipeline/detectors/tracknet_runner.py::trajectory_to_action_windows`; cfg `config.json
indexing.{wasb,tracknet}.velocity_threshold`, `dead_time_merge_gap`,
`min_action_window_duration` |
267     | Add a segmenter | subclass `@backend/pipeline/segmenters/base.py::BasePlaySegmenter`;
`segmenter_registry.register(...)` at module bottom; import in
`@backend/pipeline/segmenters/__init__.py`; if AI-based add to
`@backend/reporting/harvest.py::_AI_SEGMENTERS` |
268     | Change stitching padding | `config.json stitching.{begin_padding,end_padding}`; impl
`@backend/pipeline/stitcher/compiler.py::VideoCompiler.compile_highlights`; mirror
`@backend/tools/rally_slicer.py::BEGIN_PADDING/END_PADDING` |
269     | Add a sport | subclass `@backend/sports/base.py::SportProfile`;
`sport_registry.register(...)` in `@backend/sports/__init__.py`; define prompt+HSV in new file |
270     | Add an AI provider | subclass `@backend/providers/base.py::AIVideoProvider`;
`provider_registry.register('name', Cls)` |

```

```

271 | Add an annotation/GT tier | subclass
`@backend/annotations/base.py::AnnotationProvider`; register + import in
`@backend/annotations/__init__.py` |
272 | Add / change a report footgun rule | `@backend/reporting/spec.yaml::rules`
(declarative; ? `when.path` must match harvest's stage/metric/detail names); facts in
`@backend/reporting/harvest.py` |
273 | Tweak what the report records | `@backend/reporting/harvest.py::record_run` + stage
builders; emit point is `@backend/main.py` `finally:` ->
`@backend/reporting/__init__.py::emit_indexer_report` |
274 | Run calibration (WASB thresholds) | `@backend/eval/calibrate_wasb.py::main`
(`--trajectory --labels --fps --frame-width`); trust `cv` recall not `fit` |
275 | Distill Gemini -> local (thresholds / learned) |
`@backend/eval/calibrate_local.py::main` (thresholds) OR `@backend/eval/distill_local.py::main`
(learned classifier); both `--manifest` |
276 | Refine WASB candidates with Gemini (tuning) | `@backend/eval/gemini_refine.py::main`
(`--full-video` skips `--trajectory`; needs `GEMINI_API_KEY`) |
277 | Export tuned rally segments / clips | `@backend/eval/export_wasb_segments.py::main`
(`--slice` needs `--video`) |
278 | Export timelines (EDL/JSON/CSV) + basic stats (Value Layer slice) |
`backend/tools/export.py` (segments_to_edl/json/csv, basic_stats); called from anywhere that has
play_segments (CLI, notebook, future API) |
279 | Run a video end-to-end | `@backend/main.py::main` (`--video-path` or `--server`) or
`@run_process_and_stitch.py::main` (edit hardcoded paths) |
280 | Score DB preds vs GT (one video) | `@run_eval.py::main` |
281 | Batch score a collector manifest | `@run_phase3_eval.py::main` (`--manifest`) |
282 | Guard a regression (CI) | `@run_regression.py::main` (`--video-id`); exit 1=regress,
2=missing DB; snapshot via `--update-baseline` |
283 | Run the whole test suite | `@run_all_tests.py` (repo root); or `pytest tests/`
directly |
284 | Resume a crashed WASB pass | nothing to do ?
`@backend/pipeline/detectors/wasb_infer.py::run` auto-resumes from
`det_raw.jsonl`+`manifest.json` next to `--out`; `--fresh` forces restart |
285 | Add a validator | subclass `@backend/pipeline/validators/base.py::VideoValidator`;
`registry.register(instance)` in `@backend/pipeline/validators/__init__.py` (mind order;
producers before consumers) |
286 | Add a DB column/table | `@backend/infrastructure/database.py::Database._init_db`
migration ladder + bump `PRAGMA user_version` (else version-assert tests fail) |
287 | Train the learned segmenter | `@backend/eval/training.py::build_training_set` ->
`@backend/eval/classifier.py::LogisticRegression.fit/save` |
288 | Probe an audio signal (diagnostic) | `@backend/eval/audio/e1_probe.py::main`; detector
in `@backend/pipeline/audio/hit_detector.py` |
289
290 ---
291
292 ## 8. TEST MAP (`tests/test_*.py` -> subsystem)
293 All pure-logic / mocked-boundary unit tests; no WSL/GPU/Gemini/ffmpeg actually invoked.
`@run_all_tests.py` (repo root) is the suite runner.
294 | Test | Covers |
295 |---|---|
296 | `test_base.py` | `pipeline/segmenters/base.SegmenterRegistry` (register/get/list via
DummySegmenter) |
297 | `test_yolo_hybrid.py` | `pipeline/segmenters/yolo_hybrid.HybridYoloSegmenter`
(torchvision+ByteTrack; `_PERSON_CLASS_ID` filter; candidate link; `store_yolo_candidates=False`
skip) |
298 | `test_tracknet_hybrid.py` |
`pipeline/segmenters/tracknet_hybrid.TrackNetHybridSegmenter` (4-phase; healthcheck-abort;
`_probe_fps_width` fallback) |
299 | `test_gemini.py` | `pipeline/segmenters/gemini.GeminiPlaySegmenter` (single vs chunked
path, per-chunk offset, `shot_count`<=`shuttle_exchange_count`) |
300 | `test_tracknet_runner.py` | `pipeline/detectors/tracknet_runner` (`to_wsl_mnt_path`,
`from_indexing_cfg`, parse + `trajectory_to_action_windows` shared brain, healthcheck,
run_predict, stage) |
301 | `test_rally_gate.py` | `pipeline/detectors/rally_gate` (axis/net estimate, crossing
count, gate kills single-toss) |
302 | `test_wasb_cache.py` | `pipeline/detectors/wasb_infer` resumable cache (serialize,

```

```

contiguity heal, manifest invalidation, atomic writes, `default_cache_dir`) |
303 | `test_eval_metrics.py` | `eval.metrics` IoU/matching/score/pr_curve |
304 | `test_eval_annotations.py` | `eval.annotations` GT CSV parse/validate,
`write_scaffold` |
305 | `test_eval_harness.py` | `eval.harness` DB-pred fetch + evaluate + report_to_dict;
also `@run_eval.py::main`/`get_file_md5` (exit codes, `--scaffold`/`--json`) |
306 | `test_eval_batch.py` | `eval.batch` path resolve / stage select / aggregate
(micro/macro, zero-div) |
307 | `test_calibration.py` | `eval.calibration`
pct_improvement/kfold/select_best/cross_validate/LOVO/improvement_report |
308 | `test_calibrate_wasb.py` | `eval.calibrate_wasb`
load_golden_gts/make_predict_fn/grid/run, BASELINE |
309 | `test_calibrate_local.py` | `eval.calibrate_local` local grid + crossing gate, leave-
one-video-out |
310 | `test_export_wasb_segments.py` | `eval.export_wasb_segments` segment/comparison CSV
schemas, mmss; export->slicer reuse |
311 | `test_eval_training.py` | `eval.training` label_candidates / feature extraction /
build_training_set |
312 | `test_eval_classifier.py` | `eval.classifier.LogisticRegression`
fit/predict_proba/balanced/JSON round-trip |
313 | `test_eval_regression.py` | `eval.regression` baseline store/compare/gate; also
`@run_regression.py::main` (exit-code CI gate) |
314 | `test_distill_local.py` | `eval.distill_local` build_candidates / predict_rallies /
learned-vs-baseline |
315 | `test_gemini_refine.py` | `eval.gemini_refine` merge_overlaps / _offset_segments (pure
helpers, no API) |
316 | `test_hit_detector.py` | `pipeline/audio/hit_detector` (detect_hits, onset_envelope,
_moving_stats) + `eval/audio/e1_probe.cluster_hits` |
317 | `test_rally_slicer.py` | `tools.rally_slicer.compute_clip_windows` (pure) + label
parse |
318 | `test_video.py` | `utils.video` duration/chunk (subprocess mocked) |
319 | `test_geometry.py` | `utils.geometry` center/distance/velocity/normalised-velocity/IoU
|
320 | `test_database.py` | `infrastructure.database.Database`
schema/CRUD/migrations/`user_version`/`EXPECTED_CANDIDATE_COLUMNS` guard |
321 | `test_config.py` | `config` (load_config defaults, save_default_config, AppConfig
defaults pinned, `extra` tolerance, `.get()` dict-compat) |
322 | `test_reporting_harvest.py` | `reporting` (`emit_indexer_report`,
`harvest.record_run`); ? `importorskip("obsreport")`; all 7 footgun rules + plain-dict-OR-
AppConfig + `report.enabled=False` |
323 | `test_reporting_spec_faq.py` | `reporting/spec.yaml` <-> `README.md` cross-ref
integrity (every `faq_ref` resolves; actionable rules cite a `faq_ref`); ?
`importorskip("obsreport")` |
324 | `test_validators_base.py` | `pipeline/validators/base.ValidationRegistry`
(register/run_all/run_all_with_context -> `(results,context)`; exception captured not raised)
|
325 | `test_sport_profiles.py` | `sports` registry + BadmintonSportProfile (prompt contains
"decimal seconds"/chunk value; relevance config keys) |
326 | `test_providers.py` / `test_provider_gemini.py` | `providers` registry +
GeminiProvider lifecycle (`MAX_UPLOAD_MB` guard, non-list JSON -> [], cleanup delete) |
327 | `test_annotation_providers.py` / `test_automated_provider.py` | `annotations`
registry, FileProvider, AutomatedProvider (oracle/lineage); keystone one-scorer contract ->
`eval.metrics.score` |
328 | (no dedicated test) | individual `pipeline/validators/*` impls,
`pipeline/stitcher.compiler`, `main.py`, `run_process_and_stitch.py`,
`pipeline/detectors/wasb_runner`, `pipeline/segmenters/{wasb_hybrid,motion}`,
`eval/audio/e1_probe` (diagnostic) ?? exercised only indirectly |
329 | ? stale `.pyc` only | `tests/__pycache__/test_gemini_api.*` ? no live source exists |
330

```

5. user (2026-06-10T05:59:49.575Z)

```

426 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/main.py
148 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/config/models.py
75 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/config/loader.py

```



```

267 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/infrastructure/database.py
16 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/segmenters/__init__.py
63 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/base.py
432 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/gemini.py
229 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/segmenters/motion.py
254 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/segmenters/tracknet_hybrid.py
225 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/segmenters/wasb_hybrid.py
572 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/segmenters/yolo_hybrid.py
21 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/detectors/__init__.py
87 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/detectors/base.py
112 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/detectors/rally_gate.py
299 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/detectors/tracknet_runner.py
410 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/detectors/wasb_infer.py
150 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/detectors/wasb_runner.py
17 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/validators/__init__.py
70 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/base.py
84 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/blur.py
86 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/validators/format.py
116 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/validators/pts_continuity.py
110 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/validators/relevance.py
67 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/validators/resolution_fps.py
102 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/validators/scene_cut.py
221 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/pipeline/stitcher/compiler.py
146 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/audio/hit_detector.py
22 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/providers/__init__.py
27 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/providers/base.py
131 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/providers/gemini.py
23 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/sports/__init__.py
70 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/sports/badminton.py
18 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/sports/base.py
11 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/__init__.py
103 /c/Users/avidu/Projects/badminton-highlight-
indexer/backend/annotations/automated_provider.py
110 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/base.py
66 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/annotations/file_provider.py
71 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/reporting/__init__.py
217 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/reporting/harvest.py
1 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/__init__.py
59 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/geometry.py
33 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/hashing.py
52 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/utils/video.py
60 /c/Users/avidu/Projects/badminton-highlight-indexer/backend/results.py
5879 total

```

6. user (2026-06-10T05:59:54.146Z)

```

1 import os
2 import sys
3 import argparse
4 import json

```

```

5     import logging
6     import uvicorn
7     from dotenv import load_dotenv
8
9     load_dotenv()
10
11     # Centralized logging setup for the indexer (hot paths use logger, CLI output uses print
for direct user-facing messages).
12     logging.basicConfig(
13         level=logging.INFO,
14         format='%(asctime)s [%(levelname)s] %(name)s: %(message)s',
15         datefmt='%H:%M:%S'
16     )
17     logger = logging.getLogger(__name__)
18     from fastapi import FastAPI, HTTPException
19     from fastapi.middleware.cors import CORSMiddleware
20     from pydantic import BaseModel
21     from typing import List, Dict, Any, Optional
22
23     from fastapi.staticfiles import StaticFiles
24     from fastapi.responses import HTMLResponse
25     from backend.pipeline.validators.base import registry
26     from backend.infrastructure.database import Database
27     from backend.pipeline.segmenters.base import segmenter_registry, BasePlaySegmenter
28     from backend.pipeline.stitcher.compiler import VideoCompiler
29     from backend.utils.hashing import compute_video_id # canonical file-identity hashing
30     from backend.config import load_config as load_app_config, AppConfig # new typed config
31     from backend.results import Result, Failure # standardized error/result contract
32     from backend.reporting import emit_indexer_report
33
34     app = FastAPI(title="Sports Highlight Indexer API")
35
36     # Enable CORS for local web interface
37     app.add_middleware(
38         CORSMiddleware,
39         allow_origins=["*"],
40         allow_credentials=True,
41         allow_methods=["*"],
42         allow_headers=["*"],
43     )
44
45     # Serves static frontend files directly (now under api/ per approved structure)
46     os.makedirs("backend/api/static", exist_ok=True)
47     app.mount("/static", StaticFiles(directory="backend/api/static"), name="static")
48
49     @app.get("/", response_class=HTMLResponse)
50     def serve_index():
51         index_path = "backend/api/static/index.html"
52         if os.path.exists(index_path):
53             with open(index_path, "r", encoding="utf-8") as f:
54                 return f.read()
55         return "<h3>Sports Highlight Indexer Server is Running. Please create index.html in
backend/api/static/</h3>"
56
57     # Global variables initialized at startup
58     db_instance: Optional[Database] = None
59     global_config: Optional[AppConfig] = None # now typed (Pydantic model)
60
61     # Legacy thin wrapper for any remaining direct calls during transition.
62     # New code should import load_app_config / AppConfig from backend.config directly.
63     def load_config(config_path: str = "config.json") -> Dict[str, Any]:
64         cfg = load_app_config(config_path)
65         return cfg.model_dump(mode="json")
66
67     # --- API Models ---

```

```

68     class ProcessRequest(BaseModel):
69         video_path: str
70         player_pool: Optional[List[str]] = None
71         max_frames: Optional[int] = None
72         segmenter: Optional[str] = None
73
74     class QueryRequest(BaseModel):
75         video_id: str
76         server: Optional[str] = None
77         receiver: Optional[str] = None
78         duration_min: Optional[float] = None
79         event_type: Optional[str] = None
80
81     class CompileRequest(BaseModel):
82         video_id: str
83         segment_ids: List[int]
84         output_filename: str
85
86     # --- API Endpoints ---
87     @app.get("/api/config")
88     def get_config():
89         return global_config
90
91     @app.post("/api/process")
92     def process_video(req: ProcessRequest):
93         if not os.path.exists(req.video_path):
94             raise HTTPException(status_code=404, detail=f"Video file not found at
95 '{req.video_path}'")
96
97         video_id = compute_video_id(req.video_path)
98         was_reingest = db_instance.get_video(video_id) is not None
99
100        # 1. Run pluggable validators (with-context variant surfaces ffprobe_meta for the
report)
101        logger.info(f"Running validations for {req.video_path}...")
102        val_results, val_context = registry.run_all_with_context(req.video_path,
global_config)
103        failures = [r for r in val_results if not r.passed]
104
105        # typed AppConfig: attribute access for the default segmenter (with safe fallback)
106        default_seg = getattr(getattr(global_config, "indexing", None), "default_segmenter",
"motion")
107        seg_name = req.segmenter or default_seg
108        segmenter_actual = None
109        segmentation_ran = False
110        response: Optional[Dict[str, Any]] = None
111        try:
112            if failures:
113                db_instance.add_video(video_id, req.video_path, "failed", [r.dict() for r in
val_results])
114                response = {
115                    "video_id": video_id,
116                    "status": "failed",
117                    "message": "Video failed validation checks.",
118                    "results": [r.dict() for r in val_results],
119                }
120            return response
121
122        # 2. Run segmenter & indexer
123        print(f"[API] Video passed checks. Segmenting and indexing...")
124        db_instance.add_video(video_id, req.video_path, "processed", [r.dict() for r in
val_results])
125
126        segmenter_cls = segmenter_registry.get(seg_name)
127        if not segmenter_cls:

```

```

127         available = list(segmenter_registry.list_available().keys())
128         raise HTTPException(
129             status_code=400,
130             detail=f"Segmenter '{seg_name}' not found. Available segmenters:
{available}"
131         )
132
133     print(f"[API] Using segmenter: {seg_name}")
134     segmenter_actual = seg_name
135     segmenter = segmenter_cls(db_instance, global_config)
136     result = segmenter.process_video(video_id, req.video_path, req.player_pool,
req.max_frames)
137     segmentation_ran = True
138
139     if isinstance(result, Failure):
140         print(f"[API] Segmenter failed: {result.message}")
141         response = {
142             "video_id": video_id,
143             "status": "failed",
144             "message": f"Segmenter '{seg_name}' failed: {result.message}",
145             "results": [r.dict() for r in val_results],
146             "play_segments": [],
147         }
148         return response
149
150     segments = result.value if hasattr(result, "value") else result
151
152     response = {
153         "video_id": video_id,
154         "status": "success",
155         "message": f"Successfully indexed {len(segments)} play segments using
'{seg_name}' segmenter.",
156         "results": [r.dict() for r in val_results],
157         "play_segments": segments,
158     }
159     return response
160     finally:
161         # Always emit the per-video observability report (next to the video, or to
162         # report.output_dir). Never raises. Surface the paths in the response.
163         paths = emit_indexer_report(
164             db=db_instance, video_id=video_id, video_path=req.video_path,
config=global_config,
165             val_results=val_results, val_context=val_context,
166             segmenter_requested=req.segmenter, segmenter_actual=segmenter_actual,
167             was_reingest=was_reingest, segmentation_ran=segmentation_ran,
168         )
169         if paths and isinstance(response, dict):
170             response["reports"] = paths
171
172     @app.post("/api/query")
173     def query_play_segments(req: QueryRequest):
174         filters = {
175             "server": req.server,
176             "receiver": req.receiver,
177             "duration_min": req.duration_min,
178             "event_type": req.event_type
179         }
180         segments = db_instance.query_play_segments(req.video_id, filters)
181         return {"play_segments": segments}
182
183     @app.post("/api/compile")
184     def compile_highlights(req: CompileRequest):
185         video = db_instance.get_video(req.video_id)
186         if not video:
187             raise HTTPException(status_code=404, detail="Indexed video record not found.")

```

```

188
189     # Retrieve matching play segments from db
190     all_segments = db_instance.query_play_segments(req.video_id, {})
191     matched_segments = [s for s in all_segments if s["id"] in req.segment_ids]
192
193     if not matched_segments:
194         raise HTTPException(status_code=400, detail="No matching play segments found to
stitch.")
195
196     # Extract start/end time pairs
197     time_pairs = [(s["start_time"], s["end_time"]) for s in matched_segments]
198
199     # Run compiler
200     output_dir = getattr(getattr(global_config, "output", None), "output_dir", "output")
or "output"
201     compiler = VideoCompiler(output_dir)
202     try:
203         out_path = compiler.compile_highlights(video["filepath"], time_pairs,
req.output_filename)
204         return {"status": "success", "filepath": out_path}
205     except Exception as e:
206         raise HTTPException(status_code=500, detail=f"Highlight stitching failed:
{str(e)}")
207
208     # --- CLI Debug Utility & Orchestration ---
209     def run_cli_diagnose_clip(video_path: str, timestamp: float):
210         """CLI debugging utility to examine segment properties around a timestamp."""
211         print("-" * 60)
212         print(f"DEBUGGING VIDEO CLIP: {video_path} at {timestamp}s")
213         print("-" * 60)
214
215         cfg = load_config() # legacy wrapper still works, returns dict for now
216
217         # 1. Validation check diagnostics
218         print("[DIAGNOSTIC] Running validation framework...")
219         results = registry.run_all(video_path, cfg)
220         for r in results:
221             status_symbol = "[PASS]" if r.passed else "[FAIL]"
222             print(f" {status_symbol} {r.validator_name}: {r.message}")
223             if r.details:
224                 print(f"      Details: {r.details}")
225
226         # 2. Local segment analysis around the timestamp (sampling +/- 5 seconds)
227         print("-" * 60)
228         print("[DIAGNOSTIC] Analyzing local motion around the timestamp...")
229         import cv2
230         cap = cv2.VideoCapture(video_path)
231         if not cap.isOpened():
232             print("[FAIL] OpenCV could not open video.")
233             return
234
235         fps = cap.get(cv2.CAP_PROP_FPS)
236         total_frames = cap.get(cv2.CAP_PROP_FRAME_COUNT)
237
238         start_frame = max(0, int((timestamp - 3.0) * fps))
239         end_frame = min(int(total_frames), int((timestamp + 3.0) * fps))
240
241         # Measure frame-by-frame changes in sample region
242         cap.set(cv2.CAP_PROP_POS_FRAMES, start_frame)
243         prev_gray = None
244         motions = []
245
246         for f in range(start_frame, end_frame):
247             ret, frame = cap.read()
248             if not ret:

```

```

249         break
250         gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
251         gray = cv2.GaussianBlur(gray, (21, 21), 0)
252
253         if prev_gray is not None:
254             diff = cv2.absdiff(prev_gray, gray)
255             _, thresh = cv2.threshold(diff, 25, 255, cv2.THRESH_BINARY)
256             motion_ratio = cv2.countNonZero(thresh) / (thresh.shape[0] *
thresh.shape[1])
257             motions.append(motion_ratio)
258             prev_gray = gray
259
260         cap.release()
261
262         if motions:
263             avg_motion = sum(motions) / len(motions)
264             # Support both legacy dict (from wrapper) and future direct model
265             if hasattr(cfg, "indexing"):
266                 threshold = getattr(cfg.indexing, "motion_threshold", 0.08)
267             else:
268                 threshold = cfg.get("indexing", {}).get("motion_threshold", 0.08) # type:
ignore[attr-defined]
269             print(f" Sample Frame Count: {len(motions)}")
270             print(f" Average motion index around timestamp: {round(avg_motion, 4)}
(threshold: {threshold})")
271             if avg_motion > threshold:
272                 print(" Result: [ACTIVE PLAY] High motion detected. Segment classifies as
active play.")
273             else:
274                 print(" Result: [DEAD TIME] Idle movement. Segment classified as dead
time.")
275         else:
276             print(" [ERROR] No visual frames were extracted in the sample range.")
277
278         print("=" * 60)
279
280     def main():
281         parser = argparse.ArgumentParser(description="Sports Highlight Indexer
Orchestrator")
282         parser.add_argument("--video-path", help="Local path to the MP4 sports video file")
283         parser.add_argument("--output-dir", default="output", help="Directory where
processed clips/highlights are saved")
284         parser.add_argument("--setup", action="store_true", help="Launch dependency checker
and setup utility")
285         parser.add_argument("--debug-clip", type=float, help="Timestamp (in seconds) of a
video clip to run detailed diagnostics on")
286         parser.add_argument("--server", action="store_true", help="Run the FastAPI local API
server")
287         parser.add_argument("--port", type=int, default=8000, help="Port to run the FastAPI
server on")
288         parser.add_argument("--max-frames", type=int, help="Limit processing to the first N
sub-sampled frames for debugging")
289
290         available_segmenters = list(segmenter_registry.list_available().keys())
291         parser.add_argument("--segmenter", choices=available_segmenters, help=f"Choose rally
detection segmenter. Available: {available_segmenters}")
292         parser.add_argument("--trim-start", type=float, help="Start time in seconds for the
debug trim segment feature")
293         parser.add_argument("--trim-end", type=float, help="End time in seconds for the
debug trim segment feature")
294         parser.add_argument("--trim-output", help="Output filename for the debug trimmed
segment (saves in output-dir unless absolute path)")
295
296         args = parser.parse_args()
297

```

```

298     if args.setup:
299         # Invoke setup diagnostics
300         import subprocess
301         print("[INFO] Invoking setup.py...")
302         subprocess.run([sys.executable, "setup.py"])
303         return
304
305     if args.trim_start is not None and args.trim_end is not None:
306         if not args.video_path:
307             print("[ERROR] Please provide --video-path to extract a segment.")
308             return
309
310         # Determine output path
311         out_filename = args.trim_output or "trimmed_segment.mp4"
312         if os.path.isabs(out_filename):
313             out_path = out_filename
314             out_dir = os.path.dirname(out_path)
315             out_name = os.path.basename(out_path)
316         else:
317             out_dir = args.output_dir
318             out_path = os.path.join(out_dir, out_filename)
319             out_name = out_filename
320
321         os.makedirs(out_dir, exist_ok=True)
322         print(f"[CLI] Extracting debug segment from {args.trim_start}s to
{args.trim_end}s from {args.video_path}...")
323
324         compiler = VideoCompiler(out_dir)
325         try:
326             res_path = compiler.compile_highlights(args.video_path, [(args.trim_start,
args.trim_end)], out_name)
327             print(f"[SUCCESS] Trimmed segment saved to: {res_path}")
328         except Exception as e:
329             print(f"[ERROR] Failed to compile trimmed segment: {e}")
330         return
331
332     if args.debug_clip is not None:
333         if not args.video_path:
334             print("[ERROR] Please provide --video-path to debug a specific clip.")
335             return
336         run_cli_diagnose_clip(args.video_path, args.debug_clip)
337         return
338
339     # Initialize Global Config and DB
340     global global_config, db_instance
341     global_config = load_app_config() # returns typed AppConfig
342
343     # The CLI --output-dir overrides the configured output directory. It now lives
344     # on the typed model (OutputConfig.output_dir), so every consumer ? including
345     # /api/compile ? reads the same value instead of a write-only runtime hack.
346     global_config.output.output_dir = args.output_dir
347     os.makedirs(global_config.output.output_dir, exist_ok=True)
348
349     db_path = os.path.join(global_config.output.output_dir,
global_config.output.db_name)
350     db_instance = Database(db_path)
351
352     # CLI processing mode (no web UI needed)
353     if args.video_path and not args.server:
354         print(f"[CLI] Processing video file: {args.video_path}")
355         if not os.path.exists(args.video_path):
356             print(f"[FAIL] Video file not found: {args.video_path}")
357             return
358
359         video_id = compute_video_id(args.video_path)

```

```

360         was_reingest = db_instance.get_video(video_id) is not None
361
362         # Validate (with-context variant surfaces ffprobe_meta for the report)
363         print("[CLI] Validating...")
364         val_results, val_context = registry.run_all_with_context(args.video_path,
global_config)
365         failures = [r for r in val_results if not r.passed]
366
367         # Save video status
368         status = "failed" if failures else "processed"
369         db_instance.add_video(video_id, args.video_path, status, [r.dict() for r in
val_results])
370
371         seg_name = args.segmenter or getattr(getattr(global_config, "indexing", None),
"default_segmenter", "motion")
372         segmenter_actual = None
373         segmentation_ran = False
374         try:
375             print("\n" + "="*40)
376             print("VALIDATION SUMMARY")
377             print("="*40)
378             for r in val_results:
379                 status_symbol = "[PASS]" if r.passed else "[FAIL]"
380                 print(f"{status_symbol} {r.validator_name}: {r.message}")
381             print("="*40 + "\n")
382
383             if failures:
384                 print("[FAIL] Video failed checks. Indexing halted.")
385                 return
386
387             # Index
388             segmenter_cls = segmenter_registry.get(seg_name)
389             if not segmenter_cls:
390                 print(f"[FAIL] Segmenter '{seg_name}' not found.")
391                 return
392
393             print(f"[CLI] Indexing play segments using segmenter: {seg_name}...")
394             segmenter_actual = seg_name
395             segmenter = segmenter_cls(db_instance, global_config)
396             result = segmenter.process_video(video_id, args.video_path,
max_frames=args.max_frames)
397             segmentation_ran = True
398
399             if isinstance(result, Failure):
400                 print(f"[FAIL] Segmenter '{seg_name}' failed: {result.message}")
401                 print(f"[FAIL] Indexing halted for video.")
402                 segments = []
403             else:
404                 segments = result.value if hasattr(result, "value") else result
405                 print(f"[SUCCESS] Process complete. Indexed {len(segments)} play
segments inside database: {db_path}")
406             finally:
407                 # Always emit the per-video observability report. Never raises.
408                 paths = emit_indexer_report(
409                     db=db_instance, video_id=video_id, video_path=args.video_path,
config=global_config,
410                     val_results=val_results, val_context=val_context,
411                     segmenter_requested=args.segmenter, segmenter_actual=segmenter_actual,
412                     was_reingest=was_reingest, segmentation_ran=segmentation_ran,
413                 )
414             if paths:
415                 print(f"[REPORT] Technical report : {paths.get('technical_md')}")
416                 print(f"[REPORT] Customer summary : {paths.get('customer_md')}")
417                 print(f"[REPORT] Machine JSON      : {paths.get('json')}")
418         return

```



```

419
420     # Start Local Server Mode
421     if args.server or not args.video_path:
422         print(f"[SERVER] Starting local FastAPI server on
http://localhost:{args.port}...")
423         uvicorn.run(app, host="127.0.0.1", port=args.port)
424
425     if __name__ == "__main__":
426         main()
427

```

7. assistant (2026-06-10T06:00:21.710Z)

Now the config system.

8. user (2026-06-10T06:00:21.868Z)

```

1     """Pydantic models for the application configuration.
2
3     This module defines the canonical, validated, typed configuration for the
4     Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6     The models support graceful migration from old config.json files via extra="allow".
7     """
8
9     from __future__ import annotations
10
11     from typing import Any, Literal
12
13     from pydantic import BaseModel, Field, field_validator
14
15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0
29         max_scene_cuts_per_minute: float = 0.5
30         relevance: ValidationRelevanceConfig =
Field(default_factory=ValidationRelevanceConfig)
31
32
33     class TrackNetConfig(BaseModel):
34         wsl_distro: str = "Ubuntu"
35         repo_dir: str = "~/models/TrackNetV4"
36         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
37         conda_env: str = "TrackNetV4"
38         weights_path: str = ""
39         queue_length: int = 5
40         stage_in_wsl: bool = True
41         wsl_stage_dir: str = "~/clips"
42         timeout_sec: int = 0
43         velocity_threshold: float = 0.02
44
45
46     class IndexingConfig(BaseModel):
47         default_segmenter: str = "yolo_hybrid"

```

```

48     use_gpu: bool = True
49     motion_threshold: float = 0.08
50     min_segment_duration: float = 3.0
51     dead_time_merge_gap: float = 2.0
52     chunk_duration: float = 90.0
53     chunk_overlap: float = 10.0
54     detector_model: str = "fasterrcnn_resnet50_fpn"
55     detector_score_threshold: float = 0.5
56     yolo_velocity_threshold: float = 0.15
57     yolo_frame_skip: int = 3
58     yolo_tracker: str = "bytetrack.yaml"
59     yolo_parallel_chunks: bool = False
60     yolo_chunk_workers: int = 2
61     min_action_window_duration: float = 2.0
62     log_yolo_candidates: bool = True
63     store_yolo_candidates: bool = True
64     ai_handoff_padding: float = 2.0
65     ai_max_workers: int = 5
66
67     tracknet: TrackNetConfig = Field(default_factory=TrackNetConfig)
68
69     @field_validator("default_segmenter")
70     @classmethod
71     def segmenter_must_be_registered(cls, v: str) -> str:
72         # Registry check happens at runtime in main/segmenter selection.
73         # Here we just allow any string for forward compatibility.
74         return v
75
76
77 class OutputConfig(BaseModel):
78     default_highlights_name: str = "highlights.mp4"
79     db_name: str = "sports_indexer.db"
80     # Directory where the DB, highlight reels, and processed clips are written.
81     # The CLI's --output-dir overrides this at startup (see backend/main.py::main).
82     output_dir: str = "output"
83
84
85 class StitchingConfig(BaseModel):
86     begin_padding: float = 0.5
87     end_padding: float = 0.5
88     use_cache: bool = False
89     min_shot_count: int = 1
90
91
92 class AppConfig(BaseModel):
93     """Root configuration object.
94
95     All runtime configuration should be accessed via an instance of this class
96     (or the loader) rather than raw dict lookups.
97
98     """
99     sport: Sport = "badminton"
100     ai_provider: str = "gemini"
101     ai_model: str = "gemini-2.5-flash"
102
103     validation: ValidationConfig = Field(default_factory=ValidationConfig)
104     indexing: IndexingConfig = Field(default_factory=IndexingConfig)
105     output: OutputConfig = Field(default_factory=OutputConfig)
106     stitching: StitchingConfig = Field(default_factory=StitchingConfig)
107
108     # Allow extra keys during transition so existing config.json files do not break.
109     model_config = {
110         "extra": "allow",
111         "validate_assignment": True,
112     }

```

```

113
114     # Temporary migration shim for legacy dict-style access (config.get("section",
115     {}).get(...)).
116     # Will be removed once all call sites (validators, segmenters, etc.) move to typed
117     # attribute access / the typed config. A nested *section* is returned as a plain
dict
118     # via model_dump, so legacy chained access keeps working ? returning the Pydantic
119     # section model instead breaks it (a model has no .get), which bit every validator
on
120     # real runs. Scalar leaves are returned as-is.
121     def get(self, key: str, default: Any = None) -> Any:
122         _MISSING = object()
123         val: Any = _MISSING
124         if key in type(self).model_fields or hasattr(self, key):
125             val = getattr(self, key, _MISSING)
126         if val is _MISSING:
127             # Legacy convenience: find a leaf key inside any nested section. Sections
are
128             # discovered dynamically from the model fields, so a future top-level
section
129             # (e.g. `audio`, `jobs`) is covered automatically without editing a
hardcoded list.
130             for fname in type(self).model_fields:
131                 section = getattr(self, fname, None)
132                 if isinstance(section, BaseModel) and hasattr(section, key):
133                     val = getattr(section, key)
134                     break
135             if val is _MISSING:
136                 return default
137             if isinstance(val, BaseModel):
138                 return val.model_dump(mode="json")
139             return val
140     def __getitem__(self, key: str) -> Any:
141         val = self.get(key)
142         if val is None and key not in type(self).model_fields:
143             # Re-raise like a real dict for unknown top-level
144             raise KeyError(key)
145         return val
146
147
148     # Convenience type alias for callers that want to annotate
149     Config = AppConfig

```

9. user (2026-06-10T06:00:22.697Z)

```

1     """Configuration loader and saver.
2
3     Provides a single place to load validated AppConfig with built-in defaults
4     merged with optional file overrides (config.json).
5
6     This replaces scattered json.load + .get() patterns throughout the codebase.
7     """
8
9     from __future__ import annotations
10
11     import json
12     import logging
13     from pathlib import Path
14     from typing import Any
15
16     from .models import AppConfig
17
18     logger = logging.getLogger(__name__)
19

```

```

20     DEFAULT_CONFIG_PATH = Path("config.json")
21
22
23     def _get_builtin_defaults() -> dict[str, Any]:
24         """Return a plain dict of the current built-in defaults.
25
26         The AppConfig model is the single source of truth for defaults.
27         """
28         return AppConfig().model_dump(mode="json")
29
30
31     def load_config(path: str | Path | None = None) -> AppConfig:
32         """Load configuration with the following precedence:
33
34         1. Built-in defaults defined in the Pydantic models.
35         2. Values from the JSON file (if present and readable).
36         3. (Future) environment / CLI overrides.
37
38         Missing file or unreadable file ? returns a fully defaulted AppConfig.
39         Unknown keys in the file are tolerated (extra="allow" on the model).
40         """
41         cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
42
43         file_data: dict[str, Any] = {}
44         if cfg_path.exists():
45             try:
46                 with open(cfg_path, "r", encoding="utf-8") as f:
47                     file_data = json.load(f)
48             except Exception as exc:
49                 # Non-fatal: continue with defaults + whatever we could parse.
50                 # In production we might log here.
51                 logger.warning(f"Failed to read {cfg_path}: {exc}. Using defaults + partial
data.")
52
53         # Merge: start from defaults, overlay file data (Pydantic will validate/coerce).
54         merged = {**_get_builtin_defaults(), **file_data}
55
56         return AppConfig.model_validate(merged)
57
58
59     def save_default_config(path: str | Path | None = None) -> Path:
60         """Write a fresh default config.json.
61
62         Used by setup.py diagnostics and the --setup CLI path.
63         Overwrites if the file exists.
64         """
65         cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
66         cfg = AppConfig()
67         cfg_path.parent.mkdir(parents=True, exist_ok=True)
68         cfg_path.write_text(cfg.model_dump_json(indent=4), encoding="utf-8")
69         return cfg_path
70
71
72     # Backwards-compatible alias used by some older call sites during transition.
73     def get_default_config() -> dict[str, Any]:
74         """Return defaults as a plain dict (for transitional use only)."""
75         return _get_builtin_defaults()
76

```

10. user (2026-06-10T06:00:22.758Z)

```

1     """Lightweight Result type for standardizing success vs. failure across the pipeline.
2
3     This is the start of the Error/Result Contract standardization (low-regret item #3).
4

```

```

5      The goal is to make it impossible to confuse "no data / no rallies" with "something
actually failed"
6      (e.g. missing API key, WSL healthcheck failure, provider error, etc.).
7
8      We use simple dataclasses so callers can use isinstance checks or structural pattern
matching.
9      This is intentionally small and incremental ? we start by applying it to the segmenter
10     and provider layers (the biggest sources of silent [] returns today).
11
12     The types are generic where it makes sense. Callers can still access .value on Success
13     or .message on Failure.
14
15     This module has no dependencies on the rest of the app.
16     """
17
18     from __future__ import annotations
19
20     from dataclasses import dataclass, field
21     from typing import Generic, TypeVar, List, Optional, Any, Union
22
23     T = TypeVar("T")
24
25
26     @dataclass
27     class Success(Generic[T]):
28         """Represents a successful outcome with an optional list of non-fatal warnings."""
29
30         value: T
31         warnings: List[str] = field(default_factory=list)
32
33         def is_success(self) -> bool:
34             return True
35
36         def unwrap(self) -> T:
37             """Return the value. Only safe to call if you already know it's a Success."""
38             return self.value
39
40
41     @dataclass
42     class Failure:
43         """Represents a failure with a human-readable message and optional details."""
44
45         message: str
46         details: Optional[dict[str, Any]] = None
47         is_recoverable: bool = False # e.g. missing API key (recoverable) vs internal crash
48
49         def is_success(self) -> bool:
50             return False
51
52         def unwrap(self) -> Any:
53             """Raise to make misuse obvious in code that assumes success."""
54             raise RuntimeError(f"Attempted to unwrap a Failure: {self.message}")
55
56
57     # Convenient alias for use in type hints: Result[list[dict], Failure]
58     # Result[ValueType] means Success[ValueType] | Failure
59     # (the Failure "error" type is fixed for this contract)
60     Result = Union[Success[T], Failure]
61

```

11. assistant (2026-06-10T06:00:48.896Z)

Now the DB layer and segmenter base + twins.

12. user (2026-06-10T06:00:49.855Z)

```

1     import sqlite3
2     import json
3     import os
4     import logging
5     from typing import List, Dict, Any, Optional
6
7     logger = logging.getLogger(__name__)
8
9     class Database:
10         def __init__(self, db_path: str):
11             self.db_path = db_path
12             self._init_db()
13
14         def _get_connection(self):
15             conn = sqlite3.connect(self.db_path)
16             conn.row_factory = sqlite3.Row
17             # Enforce ON DELETE CASCADE / SET NULL ? SQLite leaves FKs off per-connection.
18             conn.execute("PRAGMA foreign_keys = ON")
19             return conn
20
21         def _init_db(self):
22             """Initializes tables for videos, play_segments, and events."""
23             with self._get_connection() as conn:
24                 cursor = conn.cursor()
25
26                 # Videos Table
27                 cursor.execute("""
28                     CREATE TABLE IF NOT EXISTS videos (
29                         id TEXT PRIMARY KEY,
30                         filepath TEXT NOT NULL,
31                         processed_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
32                         status TEXT NOT NULL,
33                         validation_results TEXT
34                     )
35                 """)
36
37                 # Play Segments Table (Extensible metadata JSON)
38                 cursor.execute("""
39                     CREATE TABLE IF NOT EXISTS play_segments (
40                         id INTEGER PRIMARY KEY AUTOINCREMENT,
41                         video_id TEXT NOT NULL,
42                         start_time REAL NOT NULL,
43                         end_time REAL NOT NULL,
44                         duration REAL NOT NULL,
45                         server TEXT,
46                         receiver TEXT,
47                         metadata TEXT,
48                         FOREIGN KEY (video_id) REFERENCES videos(id) ON DELETE CASCADE
49                     )
50                 """)
51
52                 # Events Table
53                 cursor.execute("""
54                     CREATE TABLE IF NOT EXISTS events (
55                         id INTEGER PRIMARY KEY AUTOINCREMENT,
56                         segment_id INTEGER NOT NULL,
57                         timestamp REAL NOT NULL,
58                         type TEXT NOT NULL,
59                         player TEXT,
60                         details TEXT,
61                         FOREIGN KEY (segment_id) REFERENCES play_segments(id) ON DELETE
62                         CASCADE
63                     )
64                 """)

```

```

65         # Versioned migrations live here. PRAGMA user_version is the schema
66         # contract ? bump it for every change and add an ordered `if version < N`
67         # block below. Tests assert the expected version, so accidental drift fails
68         version = cursor.execute("PRAGMA user_version").fetchone()[0]
69
70         if version < 1:
71             # v1: raw candidate detections (pre-confirmation), detector-agnostic.
72             # Common fields are columns; detector-specific metrics go in `stats`
73             # so a new segmenter reuses this table by setting its own
74             cursor.execute("""
75                 CREATE TABLE IF NOT EXISTS candidate_segments (
76                     id INTEGER PRIMARY KEY AUTOINCREMENT,
77                     video_id TEXT NOT NULL,
78                     detector TEXT NOT NULL,
79                     chunk_index INTEGER,
80                     start_time REAL NOT NULL,
81                     end_time REAL NOT NULL,
82                     duration REAL NOT NULL,
83                     confidence REAL,
84                     stats TEXT,
85                     detection_params TEXT,
86                     confirmed_segment_id INTEGER,
87                     human_label TEXT,
88                     notes TEXT,
89                     created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
90                     FOREIGN KEY (video_id) REFERENCES videos(id) ON DELETE CASCADE,
91                     FOREIGN KEY (confirmed_segment_id) REFERENCES play_segments(id)
92                 )
93             """)
94             cursor.execute("""
95                 CREATE INDEX IF NOT EXISTS idx_candidates_video_detector
96                 ON candidate_segments(video_id, detector)
97             """)
98             cursor.execute("PRAGMA user_version = 1")
99         # future: if version < 2: cursor.execute("ALTER TABLE ..."); PRAGMA
100         user_version = 2
101         conn.commit()
102
103         def add_video(self, video_id: str, filepath: str, status: str, validation_results:
104             List[Dict[str, Any]]) -> bool:
105             try:
106                 with self._get_connection() as conn:
107                     conn.cursor().execute(
108                         "INSERT OR REPLACE INTO videos (id, filepath, status,
109                         validation_results) VALUES (?, ?, ?, ?)",
110                         (video_id, filepath, status, json.dumps(validation_results))
111                     )
112                     conn.commit()
113                 return True
114             except Exception as e:
115                 logger.error(f"Failed to add video: {e}")
116                 return False
117
118         def add_play_segment(self, video_id: str, start_time: float, end_time: float,
119             server: Optional[str] = None, receiver: Optional[str] = None, metadata: Optional[Dict[str, Any]]
120             = None) -> Optional[int]:
121             try:
122                 with self._get_connection() as conn:
123                     cursor = conn.cursor()
124                     cursor.execute(

```

```

121             "INSERT INTO play_segments (video_id, start_time, end_time,
duration, server, receiver, metadata) VALUES (?, ?, ?, ?, ?, ?, ?)",
122             (video_id, start_time, end_time, end_time - start_time, server,
receiver, json.dumps(metadata or {})))
123         )
124         conn.commit()
125         return cursor.lastrowid
126     except Exception as e:
127         logger.error(f"Failed to add play segment: {e}")
128         return None
129
130     def add_event(self, segment_id: int, timestamp: float, event_type: str, player:
Optional[str] = None, details: Optional[Dict[str, Any]] = None) -> bool:
131         try:
132             with self._get_connection() as conn:
133                 conn.cursor().execute(
134                     "INSERT INTO events (segment_id, timestamp, type, player, details)
VALUES (?, ?, ?, ?, ?)",
135                     (segment_id, timestamp, event_type, player, json.dumps(details or
{})))
136                 )
137                 conn.commit()
138                 return True
139             except Exception as e:
140                 logger.error(f"Failed to add event: {e}")
141                 return False
142
143     def add_candidate_segment(self, video_id: str, detector: str, start_time: float,
end_time: float,
144                               chunk_index: Optional[int] = None, confidence:
Optional[float] = None,
145                               stats: Optional[Dict[str, Any]] = None,
146                               detection_params: Optional[Dict[str, Any]] = None) ->
Optional[int]:
147         """Persist one raw (pre-confirmation) detection. `stats`/`detection_params` are
148         detector-specific dicts stored as JSON. Returns the new candidate id, or
149         None."""
150         try:
151             with self._get_connection() as conn:
152                 cursor = conn.cursor()
153                 cursor.execute(
154                     """INSERT INTO candidate_segments
155                     (video_id, detector, chunk_index, start_time, end_time, duration,
156                     confidence, stats, detection_params)
157                     VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)""",
158                     (video_id, detector, chunk_index, start_time, end_time, end_time -
start_time,
159                     confidence, json.dumps(stats or {}), json.dumps(detection_params or
{})))
160                 )
161                 conn.commit()
162                 return cursor.lastrowid
163             except Exception as e:
164                 logger.error(f"Failed to add candidate segment: {e}")
165                 return None
166
167     def link_candidate_to_segment(self, candidate_id: int, segment_id: int) -> bool:
168         """Record that a candidate detection was confirmed as a given play_segment."""
169         try:
170             with self._get_connection() as conn:
171                 conn.cursor().execute(
172                     "UPDATE candidate_segments SET confirmed_segment_id = ? WHERE id =
?",
173                     (segment_id, candidate_id)

```



```

174         conn.commit()
175         return True
176     except Exception as e:
177         logger.error(f"Failed to link candidate {candidate_id} to segment
{segment_id}: {e}")
178         return False
179
180     def query_candidate_segments(self, video_id: str, detector: Optional[str] = None,
181                                only_unlabeled: bool = False) -> List[Dict[str, Any]]:
182         """Fetch candidate detections for the annotation / fine-tuning workflow.
183         `stats` and `detection_params` are deserialized from JSON."""
184         query = "SELECT * FROM candidate_segments WHERE video_id = ?"
185         params: List[Any] = [video_id]
186
187         if detector:
188             query += " AND detector = ?"
189             params.append(detector)
190         if only_unlabeled:
191             query += " AND human_label IS NULL"
192
193         query += " ORDER BY start_time"
194
195         candidates = []
196         with self._get_connection() as conn:
197             rows = conn.cursor().execute(query, params).fetchall()
198             for row in rows:
199                 d = dict(row)
200                 d["stats"] = json.loads(d["stats"] or "{}")
201                 d["detection_params"] = json.loads(d["detection_params"] or "{}")
202                 candidates.append(d)
203         return candidates
204
205     def get_video(self, video_id: str) -> Optional[Dict[str, Any]]:
206         with self._get_connection() as conn:
207             row = conn.cursor().execute("SELECT * FROM videos WHERE id = ?",
(video_id,)).fetchone()
208             if row:
209                 d = dict(row)
210                 d["validation_results"] = json.loads(d["validation_results"] or "[]")
211                 return d
212         return None
213
214     def query_play_segments(self, video_id: str, filters: Dict[str, Any]) ->
List[Dict[str, Any]]:
215         """
216         Dynamically queries play segments based on filters:
217         - duration_min: float
218         - server: string
219         - receiver: string
220         - event_type: string (e.g. smash, foul)
221         """
222         query = "SELECT r.* FROM play_segments r WHERE r.video_id = ?"
223         params = [video_id]
224
225         if filters.get("duration_min"):
226             query += " AND r.duration >= ?"
227             params.append(filters["duration_min"])
228
229         if filters.get("server"):
230             query += " AND r.server = ?"
231             params.append(filters["server"])
232
233         if filters.get("receiver"):
234             query += " AND r.receiver = ?"
235             params.append(filters["receiver"])

```

```

236
237         if filters.get("event_type"):
238             # Check if there is an event of a specific type inside this segment
239             query += " AND EXISTS (SELECT 1 FROM events e WHERE e.segment_id = r.id AND
e.type = ?)"
240             params.append(filters["event_type"])
241
242         segments_list = []
243         with self._get_connection() as conn:
244             cursor = conn.cursor()
245             rows = cursor.execute(query, params).fetchall()
246
247             for row in rows:
248                 r_dict = dict(row)
249                 r_id = r_dict["id"]
250                 r_dict["metadata"] = json.loads(r_dict["metadata"] or "{}")
251
252                 # Fetch associated events
253                 event_rows = cursor.execute("SELECT * FROM events WHERE segment_id = ?",
(r_id,)).fetchall()
254                 r_dict["events"] = []
255                 for er in event_rows:
256                     e_dict = dict(er)
257                     e_dict["details"] = json.loads(e_dict["details"] or "{}")
258                     r_dict["events"].append(e_dict)
259
260                 segments_list.append(r_dict)
261
262         return segments_list
263
264     def clear_video_data(self, video_id: str):
265         with self._get_connection() as conn:
266             conn.cursor().execute("DELETE FROM videos WHERE id = ?", (video_id,))
267             conn.commit()
268

```

13. user (2026-06-10T06:00:50.229Z)

```

1     from abc import ABC, abstractmethod
2     from typing import List, Dict, Any, Type, Optional
3     from backend.infrastructure.database import Database
4     from backend.results import Result # new standardized contract (Result[Value] ==
Success[Value] | Failure)
5
6     class BasePlaySegmenter(ABC):
7         def __init__(self, db: Database, config: Dict[str, Any]):
8             self.db = db
9             self.config = config
10
11         @property
12         @abstractmethod
13         def name(self) -> str:
14             """Unique identifier for this segmenter."""
15             pass
16
17         @property
18         @abstractmethod
19         def description(self) -> str:
20             """Human-readable explanation of this segmenter."""
21             pass
22
23         @abstractmethod
24         def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None, max_frames: int = None) -> "Result[List[Dict[str, Any]]]":
25             """

```

```

26         Processes a video to detect play segments, populates SQLite database,
27         and returns a Result.
28
29         Use isinstance(result, Failure) (or result.is_success()) to distinguish
30         "no segments found" from "something went wrong" (missing key, healthcheck
31         failure, etc.). This is the start of the standardized error/result contract.
32         """
33         pass
34
35     class SegmenterRegistry:
36         def __init__(self):
37             self._segmenters: Dict[str, Type[BasePlaySegmenter]] = {}
38
39         def register(self, segmenter_cls: Type[BasePlaySegmenter]):
40             """Registers a new segmenter class under its name property."""
41             try:
42                 temp_inst = segmenter_cls(None, {})
43                 name = temp_inst.name
44             except Exception as e:
45                 raise ValueError(f"Failed to resolve name property from segmenter class
{segmenter_cls.__name__}: {e}")
46             self._segmenters[name] = segmenter_cls
47
48         def get(self, name: str) -> Optional[Type[BasePlaySegmenter]]:
49             return self._segmenters.get(name)
50
51         def list_available(self) -> Dict[str, str]:
52             """Returns a dict of registered segmenter names and their descriptions."""
53             res = {}
54             for name, cls in self._segmenters.items():
55                 try:
56                     temp_inst = cls(None, {})
57                     desc = temp_inst.description
58                 except Exception:
59                     desc = "No description available."
60                 res[name] = desc
61             return res
62
63     segmenter_registry = SegmenterRegistry()
64

```

14. assistant (2026-06-10T06:01:09.774Z)

Now the five segmenters.

15. user (2026-06-10T06:01:10.531Z)

```

1     """WASB-based hybrid play segmenter.
2
3     Identical three-phase shape to TrackNetHybridSegmenter, but Phase 2 uses **WASB**
4     (nttcom/WASB-SBDT, MIT) shuttle-trajectory detection via WSL2
5     (backend/detectors/wasb_runner.py) instead of TrackNetV4. WASB ships pretrained
6     badminton weights and runs today, so this is the live shuttle-motion substrate.
7     Phase 3 (AI semantic boundaries) and Phase 4 (DB) are shared with the other hybrids.
8
9     For the per-video *calibration* goal we mainly use Phase 2 (trajectory -> candidate
10    windows); the windowing thresholds (velocity_thresh / merge_gap / min_window_duration)
11    are what get tuned against golden rally labels.
12    """
13    import os
14    import time
15    import logging
16    import concurrent.futures
17    from typing import List, Dict, Any
18

```

```

19     import cv2
20
21     from backend.pipeline.segmenters.base import BasePlaySegmenter, segmenter_registry
22     from backend.utils.video import get_video_duration, extract_video_chunk
23     from backend.sports import sport_registry
24     from backend.providers import provider_registry
25     from backend.pipeline.detectors.wasb_runner import WasbRunner, WasbConfig
26     from backend.pipeline.detectors.base import DetectorRunner # ABC for the de-risked
detector seam
27
28     logger = logging.getLogger(__name__)
29
30
31     def _fmt_ts(seconds: float) -> str:
32         seconds = max(0, int(seconds))
33         h, rem = divmod(seconds, 3600)
34         m, s = divmod(rem, 60)
35         return f"{h:02d}:{m:02d}:{s:02d}"
36
37
38     class WasbHybridSegmenter(BasePlaySegmenter):
39         @property
40         def name(self) -> str:
41             return "wasb_hybrid"
42
43         @property
44         def description(self) -> str:
45             return ("Two-stage pipeline: WASB shuttle-trajectory detection (WSL2, pretrained
badminton "
46                     "weights) for precise candidate rallies + AI provider for semantic
boundaries.")
47
48         @staticmethod
49         def _probe_fps_width(video_path: str) -> tuple:
50             cap = cv2.VideoCapture(video_path)
51             fps = cap.get(cv2.CAP_PROP_FPS) if cap.isOpened() else 0.0
52             width = cap.get(cv2.CAP_PROP_FRAME_WIDTH) if cap.isOpened() else 0.0
53             cap.release()
54             return (fps if fps > 0 else 30.0, width if width > 0 else 1280.0)
55
56         def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None,
57                         max_frames: int = None) -> List[Dict[str, Any]]:
58             sport_name = self.config.get("sport", "badminton")
59             try:
60                 sport_profile = sport_registry.get(sport_name)
61             except KeyError as e:
62                 logger.error(str(e))
63                 return []
64
65             idx_cfg = self.config.get("indexing", {})
66             provider_name = self.config.get("ai_provider", idx_cfg.get("ai_provider",
"gemini"))
67             model_name = self.config.get("ai_model", idx_cfg.get("ai_model",
"gemini-2.5-flash"))
68
69             api_key_env_var = f"{provider_name.upper()}_API_KEY"
70             api_key = os.environ.get(api_key_env_var)
71             if not api_key:
72                 logger.error(f"{api_key_env_var} not set; cannot run the {provider_name}
handoff.")
73                 return []
74             try:
75                 provider = provider_registry.get(provider_name, api_key=api_key,
model_name=model_name)

```

```

76         except KeyError as e:
77             logger.error(str(e))
78             return []
79
80         video_duration = get_video_duration(video_path)
81         if video_duration <= 0:
82             logger.error("Invalid video duration.")
83             return []
84         fps, frame_width = self._probe_fps_width(video_path)
85
86         wasb_cfg = WasbConfig.from_indexing_cfg(idx_cfg)
87         runner: DetectorRunner = WasbRunner(wasb_cfg)
88
89         w = idx_cfg.get("wasb", {})
90         velocity_thresh = w.get("velocity_threshold", 0.02)
91         merge_gap = idx_cfg.get("dead_time_merge_gap", 2.0)
92         min_window_duration = idx_cfg.get("min_action_window_duration", 2.0)
93         handoff_padding = idx_cfg.get("ai_handoff_padding", 2.0)
94         ai_max_workers = idx_cfg.get("ai_max_workers", 5)
95         log_candidates = idx_cfg.get("log_yolo_candidates", True)
96         store_candidates = idx_cfg.get("store_yolo_candidates", True)
97
98         detection_params = {
99             "detector": "wasb_hybrid",
100             "velocity_thresh": velocity_thresh,
101             "merge_gap": merge_gap,
102             "min_action_window_duration": min_window_duration,
103             "weights_path": wasb_cfg.weights_path,
104         }
105
106         ok, msg = runner.healthcheck()
107         if not ok:
108             logger.error("WASB WSL environment not ready: %s", msg)
109             return []
110         logger.info("WASB WSL env OK: %s", msg)
111
112         # --- Phase 2: shuttle trajectory -> candidate rally windows ---
113         t_start = time.time()
114         out_dir = os.path.join("output", "wasb")
115         csv_path = runner.run_predict(video_path, out_dir)
116         if not csv_path:
117             logger.error("WASB produced no trajectory; aborting.")
118             return []
119
120         points = runner.parse_trajectory_csv(csv_path)
121         logger.info("Parsed %d trajectory points (%d visible).",
122                     len(points), sum(1 for p in points if p.visible))
123
124         all_candidate_windows = runner.trajectory_to_action_windows(
125             points, fps=fps, frame_width=frame_width, chunk_start=0.0,
126             velocity_thresh=velocity_thresh, merge_gap=merge_gap,
127             min_window_duration=min_window_duration, chunk_index=0,
128         )
129         logger.info("Phase 2 (WASB) completed in %.1fs. Candidate windows: %d",
130                     time.time() - t_start, len(all_candidate_windows))
131
132         if log_candidates:
133             for cw in all_candidate_windows:
134                 logger.info(f"[WASB rally] {_fmt_ts(cw['start'])}-{_fmt_ts(cw['end'])} "
135                             f"({_fmt_ts(cw['end'] - cw['start']:.1f)s) | "
136                             f"peak_v={cw['peak_velocity']:.3f} "
137                             f"frames={cw['active_frame_count']} "
138                             f"mean_v={cw['mean_velocity']:.3f}")
139
140         candidate_ids = [None] * len(all_candidate_windows)

```

```

139         if store_candidates:
140             for i, cw in enumerate(all_candidate_windows):
141                 stats = {
142                     "peak_velocity": cw["peak_velocity"], "mean_velocity":
143                     "active_frame_count": cw["active_frame_count"], "track_id_count":
144                     cw["track_id_count"],
145                     }
146                 candidate_ids[i] = self.db.add_candidate_segment(
147                     video_id, detector="wasb_hybrid", start_time=cw["start"],
148                     end_time=cw["end"],
149                     chunk_index=cw["chunk_index"], confidence=min(1.0,
150                     cw["peak_velocity"]),
151                     stats=stats, detection_params=detection_params,
152                     )
153
154         if not all_candidate_windows:
155             return []
156
157         # --- Phase 3: AI provider handoff (shared pattern) ---
158         handoff_dir = os.path.join("output", f"chunks_{provider_name}_handoff")
159         os.makedirs(handoff_dir, exist_ok=True)
160         logger.info("Phase 3: %s validation on %d candidate windows...",
161                     provider_name, len(all_candidate_windows))
162
163         def process_candidate_window(wi: int, window: Dict[str, Any]) -> List[dict]:
164             w_start = max(0, window["start"] - handoff_padding)
165             w_end = min(video_duration, window["end"] + handoff_padding)
166             w_dur = w_end - w_start
167             w_path = os.path.join(handoff_dir, f"handoff_{video_id}_{wi}.mp4")
168             if not extract_video_chunk(video_path, w_start, w_dur, w_path,
169             reencode=True):
170                 return []
171             prompt = sport_profile.get_prompt(chunk_duration=w_dur)
172             segments, _ = provider.analyze_video(w_path, prompt, chunk_duration=w_dur,
173                                                 label=f"Handoff {wi+1}")
174             valid = []
175             for seg in segments:
176                 try:
177                     seg["start_time"] = float(seg["start_time"]) + w_start
178                     seg["end_time"] = float(seg["end_time"]) + w_start
179                     seg["_candidate_id"] = candidate_ids[wi]
180                     valid.append(seg)
181                 except (ValueError, TypeError, KeyError) as e:
182                     logger.warning("Dropping segment with unparseable timestamps: %s ?
183 %s", seg, e)
184
185             try:
186                 os.remove(w_path)
187             except OSError:
188                 pass
189             return valid
190
191         ai_segments: List[dict] = []
192         with concurrent.futures.ThreadPoolExecutor(max_workers=ai_max_workers) as
193         executor:
194             futures = [executor.submit(process_candidate_window, wi, cw)
195                         for wi, cw in enumerate(all_candidate_windows)]
196             for future in futures:
197                 ai_segments.extend(future.result())
198
199         try:
200             os.rmdir(handoff_dir)
201         except OSError:
202             pass
203         logger.info("Phase 3 completed. AI confirmed %d semantic play segments.",
204                     len(ai_segments))

```

```

196
197     # --- Phase 4: DB population (shared pattern) ---
198     segments_data = []
199     for segment in ai_segments:
200         try:
201             start = float(segment["start_time"])
202             end = float(segment["end_time"])
203             dur = end - start
204             shot_count = segment.get("shuttle_exchange_count")
205             try:
206                 shot_count = int(shot_count)
207             except (ValueError, TypeError):
208                 shot_count = max(3, int(dur / 0.8))
209             metadata = {"shot_count": shot_count,
210                        "shot_count_confidence":
211                        segment.get("shot_count_confidence", "Unknown"),
212                        "detector": f"wasb-hybrid-{model_name}"}
213             segment_id = self.db.add_play_segment(video_id, start, end, None, None,
214             metadata)
215             if segment_id:
216                 candidate_id = segment.get("_candidate_id")
217                 if candidate_id is not None:
218                     self.db.link_candidate_to_segment(candidate_id, segment_id)
219                 segments_data.append({"id": segment_id, "start_time": start,
220                                     "end_time": end,
221                                     "duration": dur, "metadata": metadata})
222             except Exception as e:
223                 logger.warning("Skipping invalid segment struct: %s", e)
224
225         return segments_data
226
227     segmenter_registry.register(WasbHybridSegmenter)
228

```

16. user (2026-06-10T06:01:11.218Z)

```

1     """TrackNet-based hybrid play segmenter.
2
3     Same three-phase shape as HybridYoloSegmenter, but Phase 2 (candidate
4     generation) uses a TrackNetV4 *shuttle* trajectory instead of YOLO *person*
5     motion. TrackNet runs in WSL2 (see backend/detectors/tracknet_runner.py and
6     docs/DECISIONS.md ADR-001); the shuttle trajectory is converted into candidate
7     rally windows, which are then handed to the AI provider for semantic boundary
8     confirmation (Phase 3) and written to the DB (Phase 4).
9
10    Rationale: a tiny, fast, motion-blurred shuttle is exactly what general
11    detectors miss; a purpose-built trajectory model localises it far more
12    precisely, giving cleaner candidate rallies than person-velocity heuristics.
13    """
14
15    import os
16    import time
17    import logging
18    import concurrent.futures
19    from typing import List, Dict, Any
20
21    import cv2
22
23    from backend.pipeline.segmenters.base import BasePlaySegmenter, segmenter_registry
24    from backend.utils.video import get_video_duration, extract_video_chunk
25    from backend.sports import sport_registry
26    from backend.providers import provider_registry
27    from backend.pipeline.detectors.tracknet_runner import TrackNetRunner, TrackNetConfig
28    from backend.pipeline.detectors.base import DetectorRunner # ABC for the de-risked

```

```

detector seam
29
30     logger = logging.getLogger(__name__)
31
32
33     def _fmt_ts(seconds: float) -> str:
34         seconds = max(0, int(seconds))
35         h, rem = divmod(seconds, 3600)
36         m, s = divmod(rem, 60)
37         return f"{h:02d}:{m:02d}:{s:02d}"
38
39
40     class TrackNetHybridSegmenter(BasePlaySegmenter):
41         @property
42         def name(self) -> str:
43             return "tracknet_hybrid"
44
45         @property
46         def description(self) -> str:
47             return ("Two-stage pipeline: TrackNetV4 shuttle-trajectory detection (WSL2) for
precise "
48                     "candidate rallies + AI provider for high-precision semantic
boundaries.")
49
50         @staticmethod
51         def _probe_fps_width(video_path: str) -> tuple:
52             """Return (fps, frame_width) for a video, with sensible fallbacks."""
53             cap = cv2.VideoCapture(video_path)
54             fps = cap.get(cv2.CAP_PROP_FPS) if cap.isOpened() else 0.0
55             width = cap.get(cv2.CAP_PROP_FRAME_WIDTH) if cap.isOpened() else 0.0
56             cap.release()
57             return (fps if fps > 0 else 30.0, width if width > 0 else 1280.0)
58
59         def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None,
60                         max_frames: int = None) -> List[Dict[str, Any]]:
61             # --- Sport profile + AI provider wiring (identical to the other segmenters) ---
62             sport_name = self.config.get("sport", "badminton")
63             try:
64                 sport_profile = sport_registry.get(sport_name)
65             except KeyError as e:
66                 logger.error(str(e))
67             return []
68
69             idx_cfg = self.config.get("indexing", {})
70             provider_name = self.config.get("ai_provider", idx_cfg.get("ai_provider",
"gemini"))
71             model_name = self.config.get("ai_model", idx_cfg.get("ai_model",
"gemini-2.5-flash"))
72
73             api_key_env_var = f"{provider_name.upper()}_API_KEY"
74             api_key = os.environ.get(api_key_env_var)
75             if not api_key:
76                 logger.error(
77                     f"{api_key_env_var} environment variable is not set! "
78                     f"To run the {provider_name} handoff, set your API key. "
79                     f"In PowerShell: $env:{api_key_env_var}=\"your_api_key_here\""
80                 )
81             return []
82             try:
83                 provider = provider_registry.get(provider_name, api_key=api_key,
model_name=model_name)
84             except KeyError as e:
85                 logger.error(str(e))
86             return []

```



```

87
88     video_duration = get_video_duration(video_path)
89     if video_duration <= 0:
90         logger.error("Invalid video duration.")
91         return []
92     fps, frame_width = self._probe_fps_width(video_path)
93
94     # --- TrackNet runner config + thresholds ---
95     tn_cfg = TrackNetConfig.from_indexing_cfg(idx_cfg)
96     runner: DetectorRunner = TrackNetRunner(tn_cfg)
97
98     tn = idx_cfg.get("tracknet", {})
99     velocity_thresh = tn.get("velocity_threshold", 0.02) # normalised shuttle
displacement/sec
100     merge_gap = idx_cfg.get("dead_time_merge_gap", 2.0)
101     min_window_duration = idx_cfg.get("min_action_window_duration", 2.0)
102     handoff_padding = idx_cfg.get("ai_handoff_padding", 2.0)
103     ai_max_workers = idx_cfg.get("ai_max_workers", 5)
104     log_candidates = idx_cfg.get("log_yolo_candidates", True)
105     store_candidates = idx_cfg.get("store_yolo_candidates", True)
106
107     detection_params = {
108         "detector": "tracknet_hybrid",
109         "velocity_thresh": velocity_thresh,
110         "merge_gap": merge_gap,
111         "min_action_window_duration": min_window_duration,
112         "weights_path": tn_cfg.weights_path,
113         "queue_length": tn_cfg.queue_length,
114     }
115
116     # --- Phase 1: env healthcheck (fail fast with a clear message) ---
117     ok, msg = runner.healthcheck()
118     if not ok:
119         logger.error("TrackNet WSL environment not ready: %s", msg)
120         return []
121     logger.info("TrackNet WSL env OK: %s", msg)
122
123     # --- Phase 2: shuttle trajectory -> candidate rally windows ---
124     # TrackNet processes the whole video in one pass (it carries its own
125     # frame buffering), so unlike the YOLO segmenter we don't pre-chunk here.
126     t_start = time.time()
127     tn_out_dir = os.path.join("output", "tracknet")
128     csv_path = runner.run_predict(video_path, tn_out_dir)
129     if not csv_path:
130         logger.error("TrackNet produced no trajectory; aborting.")
131         return []
132
133     points = runner.parse_trajectory_csv(csv_path)
134     logger.info("Parsed %d trajectory points (%d visible).",
135                 len(points), sum(1 for p in points if p.visible))
136
137     all_candidate_windows = runner.trajectory_to_action_windows(
138         points, fps=fps, frame_width=frame_width, chunk_start=0.0,
139         velocity_thresh=velocity_thresh, merge_gap=merge_gap,
140         min_window_duration=min_window_duration, chunk_index=0,
141     )
142     logger.info("Phase 2 (TrackNet) completed in %.1fs. Candidate windows: %d",
143                 time.time() - t_start, len(all_candidate_windows))
144
145     if log_candidates:
146         for w in all_candidate_windows:
147             logger.info(
148                 f"[TrackNet rally] {_fmt_ts(w['start'])}-{_fmt_ts(w['end'])} "
149                 f"({w['end'] - w['start']:.1f}s) | frames={w['active_frame_count']}"
150             )

```

```

150         f"peak_v={w['peak_velocity']:.3f} mean_v={w['mean_velocity']:.3f}"
151     )
152
153     # Persist raw candidates for the fine-tuning dataset (same table as YOLO).
154     candidate_ids = [None] * len(all_candidate_windows)
155     if store_candidates:
156         for i, w in enumerate(all_candidate_windows):
157             stats = {
158                 "peak_velocity": w["peak_velocity"],
159                 "mean_velocity": w["mean_velocity"],
160                 "active_frame_count": w["active_frame_count"],
161                 "track_id_count": w["track_id_count"],
162             }
163             confidence = min(1.0, w["peak_velocity"])
164             candidate_ids[i] = self.db.add_candidate_segment(
165                 video_id, detector="tracknet_hybrid",
166                 start_time=w["start"], end_time=w["end"],
167                 chunk_index=w["chunk_index"], confidence=confidence,
168                 stats=stats, detection_params=detection_params,
169             )
170
171     if not all_candidate_windows:
172         return []
173
174     # --- Phase 3: AI provider handoff (identical pattern to HybridYoloSegmenter)
175     ---
176     handoff_dir = os.path.join("output", f"chunks_{provider_name}_handoff")
177     os.makedirs(handoff_dir, exist_ok=True)
178     logger.info("Phase 3: %s validation on %d candidate windows...",
179                 provider_name, len(all_candidate_windows))
180
181     def process_candidate_window(wi: int, window: Dict[str, Any]) -> List[dict]:
182         w_start = max(0, window["start"] - handoff_padding)
183         w_end = min(video_duration, window["end"] + handoff_padding)
184         w_dur = w_end - w_start
185         w_path = os.path.join(handoff_dir, f"handoff_{video_id}_{wi}.mp4")
186         if not extract_video_chunk(video_path, w_start, w_dur, w_path,
187                                   reencode=True):
188             return []
189         prompt = sport_profile.get_prompt(chunk_duration=w_dur)
190         segments, _ = provider.analyze_video(w_path, prompt, chunk_duration=w_dur,
191                                              label=f"Handoff {wi+1}")
192
193         valid = []
194         for seg in segments:
195             try:
196                 seg["start_time"] = float(seg["start_time"]) + w_start
197                 seg["end_time"] = float(seg["end_time"]) + w_start
198                 seg["_candidate_id"] = candidate_ids[wi]
199                 valid.append(seg)
200             except (ValueError, TypeError, KeyError) as e:
201                 logger.warning("Dropping segment with unparseable timestamps: %s ?
202                                %s", seg, e)
203
204         try:
205             os.remove(w_path)
206         except OSError:
207             pass
208         return valid
209
210     ai_segments: List[dict] = []
211     with concurrent.futures.ThreadPoolExecutor(max_workers=ai_max_workers) as
212     executor:
213         futures = [executor.submit(process_candidate_window, wi, w)
214                     for wi, w in enumerate(all_candidate_windows)]
215         for future in futures:
216             ai_segments.extend(future.result())

```

```

211
212         try:
213             os.rmdir(handoff_dir)
214         except OSError:
215             pass
216         logger.info("Phase 3 completed. AI confirmed %d semantic play segments.",
len(ai_segments))
217
218         # --- Phase 4: DB population (identical pattern to HybridYoloSegmenter) ---
219         segments_data = []
220         for segment in ai_segments:
221             try:
222                 start = float(segment["start_time"])
223                 end = float(segment["end_time"])
224                 dur = end - start
225                 shot_count = segment.get("shuttle_exchange_count")
226                 if shot_count is None:
227                     shot_count = max(3, int(dur / 0.8))
228                 else:
229                     try:
230                         shot_count = int(shot_count)
231                     except (ValueError, TypeError):
232                         shot_count = max(3, int(dur / 0.8))
233                 metadata = {
234                     "shot_count": shot_count,
235                     "shot_count_confidence": segment.get("shot_count_confidence",
"Unknown"),
236                     "detector": f"tracknet-hybrid-{model_name}",
237                 }
238                 segment_id = self.db.add_play_segment(video_id, start, end, None, None,
metadata)
239                 if segment_id:
240                     candidate_id = segment.get("_candidate_id")
241                     if candidate_id is not None:
242                         self.db.link_candidate_to_segment(candidate_id, segment_id)
243                     segments_data.append({
244                         "id": segment_id, "start_time": start, "end_time": end,
245                         "duration": dur, "metadata": metadata,
246                     })
247             except Exception as e:
248                 logger.warning("Skipping invalid segment struct: %s", e)
249
250         return segments_data
251
252
253     # Register the segmenter
254     segmenter_registry.register(TrackNetHybridSegmenter)
255

```

17. user (2026-06-10T06:01:42.675Z)

```

1     """Two-stage rally segmenter: torchvision person detection + ByteTrack tracking
2     (Stage 1, local) feeding an AI provider for semantic boundaries (Stage 2).
3
4     Stage 1 originally used the AGPL-licensed ultralytics YOLO. To keep the hosted
5     service free of copyleft obligations (see docs/MONETIZATION_AUDIT.md and ADR-005),
6     detection now uses torchvision's permissively-licensed (BSD) detectors and the
7     association/tracking that YOLO.track() used to bundle in is provided by
8     supervision's MIT-licensed ByteTrack.
9
10    The registry key stays "yolo_hybrid" for back-compat with existing configs and the
11    DB `detector` tags, even though no YOLO is involved anymore.
12    """
13    import os
14    import time

```

```

15     import logging
16     import cv2
17     import torch
18     import concurrent.futures
19     from typing import List, Dict, Any, Tuple
20
21     from backend.pipeline.segmenters.base import BasePlaySegmenter, segmenter_registry
22     from backend.utils.video import get_video_duration, extract_video_chunk
23     from backend.utils.geometry import get_velocity, get_velocity_normalised
24     from backend.sports import sport_registry
25     from backend.providers import provider_registry
26
27     logger = logging.getLogger(__name__)
28
29     # torchvision detection models are trained on COCO where the "person" category is
30     # class 1 (index 0 is the implicit background). NB: ultralytics YOLO used class 0 ?
31     # this off-by-one is the classic gotcha when swapping detectors. Keep it named.
32     _PERSON_CLASS_ID = 1
33
34     # Permissively-licensed torchvision detectors selectable via config
35     # `indexing.detector_model`. All are BSD-licensed (code + COCO weights).
36     _DETECTOR_FACTORIES = {
37         "fasterrcnn_resnet50_fpn": "fasterrcnn_resnet50_fpn",
38         "fasterrcnn_mobilenet_v3_large_fpn": "fasterrcnn_mobilenet_v3_large_fpn",
39         "retinanet_resnet50_fpn": "retinanet_resnet50_fpn",
40     }
41     _DEFAULT_DETECTOR = "fasterrcnn_resnet50_fpn"
42
43
44     def _fmt_ts(seconds: float) -> str:
45         """Format seconds as HH:MM:SS for cross-verifying a timestamp against the video."""
46         seconds = max(0, int(seconds))
47         h, rem = divmod(seconds, 3600)
48         m, s = divmod(rem, 60)
49         return f"{h:02d}:{m:02d}:{s:02d}"
50
51
52     class HybridYoloSegmenter(BasePlaySegmenter):
53         def __init__(self, db, config):
54             super().__init__(db, config)
55             self.model = None
56             self.device = None
57
58         @property
59         def name(self) -> str:
60             return "yolo_hybrid"
61
62         @property
63         def description(self) -> str:
64             return ("Two-stage pipeline: torchvision person detection + ByteTrack for fast "
65                 "candidate filtering, then an AI provider for high-precision semantic
66                 boundaries.")
67
68         def _lazy_load_model(self):
69             if self.model is not None:
70                 return
71
72             # Imported lazily so the module imports without torchvision present and so
73             # the test suite (which pre-sets self.model) never pulls in model weights.
74             from torchvision.models import detection as tv_detection
75
76             idx_cfg = self.config.get("indexing", {})
77             use_gpu = idx_cfg.get("use_gpu", True)
78             self.device = "cuda" if use_gpu and torch.cuda.is_available() else "cpu"
79             if use_gpu and not torch.cuda.is_available():

```

```

79         logger.warning("use_gpu is True but CUDA is not available. Falling back to
CPU.")
80
81         model_name = idx_cfg.get("detector_model", _DEFAULT_DETECTOR)
82         if model_name not in _DETECTOR_FACTORIES:
83             logger.warning(f"Unknown detector_model '{model_name}'. Falling back to
{_DEFAULT_DETECTOR}.")
84             model_name = _DEFAULT_DETECTOR
85
86         builder = getattr(tv_detection, _DETECTOR_FACTORIES[model_name])
87         logger.info(f"Loading torchvision detector '{model_name}' on {self.device}...")
88         # weights="DEFAULT" pulls the best available COCO-pretrained checkpoint.
89         self.model = builder(weights="DEFAULT")
90         self.model.eval()
91         self.model.to(self.device)
92
93     def _make_tracker(self, fps: float):
94         """Create a fresh ByteTrack tracker for one chunk (track IDs are chunk-local).
95
96         Imported lazily so tests can mock this method without supervision installed.
97         """
98         import supervision as sv
99         frame_rate = max(1, int(round(fps)))
100         return sv.ByteTrack(frame_rate=frame_rate)
101
102     def _detect_and_track(self, frame, tracker, score_thresh: float) -> List[Tuple[int,
Tuple[float, float, float, float]]]:
103         """Detect persons in a frame and assign persistent track IDs.
104
105         Returns a list of (track_id, bbox) where bbox is (x1, y1, x2, y2) in pixels.
106         This replaces the old ultralytics ``model.track(...)`` call, which bundled
107         detection AND association: torchvision does the detection, supervision's
108         ByteTrack does the association.
109         """
110         import numpy as np
111         import supervision as sv
112
113         # cv2 frames are BGR uint8 HxWxC; torchvision wants RGB float [0,1] CxHxW.
114         rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
115         tensor = torch.from_numpy(rgb).permute(2, 0,
1).float().div(255.0).to(self.device)
116         with torch.no_grad():
117             outputs = self.model([tensor])
118
119         out = outputs[0]
120         boxes = out["boxes"].detach().cpu().numpy()
121         labels = out["labels"].detach().cpu().numpy()
122         scores = out["scores"].detach().cpu().numpy()
123
124         # Keep only confident person detections.
125         mask = (labels == _PERSON_CLASS_ID) & (scores >= score_thresh)
126         if not mask.any():
127             return []
128
129         # ByteTrack requires confidence to be populated (it filters on it internally).
130         dets = sv.Detections(xyxy=boxes[mask], confidence=scores[mask],
class_id=labels[mask])
131         tracked = tracker.update_with_detections(dets)
132
133         result: List[Tuple[int, Tuple[float, float, float, float]]] = []
134         for i in range(len(tracked)):
135             tid = tracked.tracker_id[i]
136             if tid is None:
137                 continue
138             x1, y1, x2, y2 = tracked.xyxy[i]

```

```

139         result.append((int(tid), (float(x1), float(y1), float(x2), float(y2))))
140     return result
141
142     def _extract_action_windows_from_chunk(self, chunk_path: str, chunk_start: float,
143         velocity_thresh: float, merge_gap: float,
144         frame_skip: int = 3,
145         tracker_config: str = "bytetrack.yaml",
146         chunk_index: int = None) -> List[Dict[str,
Any]]:
147         """
148         Runs person detection + tracking on a video chunk and extracts high-motion
149         action windows. Returns a list of dicts in the full video's timeframe, each
with:
150             start, end (seconds), active_frame_count, peak_velocity, mean_velocity,
151             track_id_count, chunk_index.
152
153         Args:
154             frame_skip: Process every Nth frame (1 = every frame). Higher values
155                 dramatically reduce inference cost with minimal quality loss.
156             tracker_config: Retained for back-compat (was the YOLO tracker yaml); the
157                 tracker is now supervision ByteTrack and this is unused.
158             chunk_index: Index of the source chunk (recorded on each window for
traceability).
159         """
160         cap = cv2.VideoCapture(chunk_path)
161         if not cap.isOpened():
162             logger.error(f"Could not open {chunk_path}")
163             return []
164
165         fps = cap.get(cv2.CAP_PROP_FPS)
166         if fps <= 0:
167             fps = 30.0
168
169         frame_width = cap.get(cv2.CAP_PROP_FRAME_WIDTH)
170         if frame_width <= 0:
171             frame_width = 1280.0 # Sensible fallback
172
173         # Effective sampling rate after frame-skipping
174         effective_fps = fps / max(frame_skip, 1)
175
176         # A fresh tracker per chunk ? track IDs only need to be consistent within a
177         # chunk (windows are computed per chunk, then merged across chunks by time).
178         tracker = self._make_tracker(effective_fps)
179         score_thresh = self.config.get("indexing", {}).get("detector_score_threshold",
0.5)
180
181         # Per active frame we keep (frame_idx, peak_velocity, set_of_active_track_ids)
182         # so windows can be enriched with motion stats for logging + fine-tuning.
183         active_frames = []
184         frame_idx = 0
185         track_history = {}
186
187         while True:
188             ret, frame = cap.read()
189             if not ret:
190                 break
191
192             # Skip frames for performance ? at 30fps, every 3rd frame still
193             # gives ~10 samples/second, plenty for human-scale velocity.
194             if frame_idx % max(frame_skip, 1) != 0:
195                 frame_idx += 1
196                 continue
197
198             # Detect persons + assign persistent track IDs.
199             detections = self._detect_and_track(frame, tracker, score_thresh)

```

```

200
201     frame_peak_velocity = 0.0
202     frame_active_ids = set()
203     current_ids = set()
204
205     for track_id, bbox in detections:
206         current_ids.add(track_id)
207
208         if track_id in track_history:
209             prev_bbox = track_history[track_id]
210             velocity = get_velocity_normalised(prev_bbox, bbox, effective_fps,
frame_width)
211
212             if velocity > velocity_thresh:
213                 frame_peak_velocity = max(frame_peak_velocity, velocity)
214                 frame_active_ids.add(track_id)
215
216             track_history[track_id] = bbox
217
218             # Prune stale IDs that are no longer visible ? prevents false
219             # velocity spikes when the tracker reassigns an old ID.
220             track_history = {k: v for k, v in track_history.items() if k in current_ids}
221
222             if frame_active_ids:
223                 active_frames.append((frame_idx, frame_peak_velocity, frame_active_ids))
224
225             frame_idx += 1
226
227     cap.release()
228
229     # Group active frames into continuous windows
230     if not active_frames:
231         return []
232
233     max_gap_frames = int(merge_gap * fps)
234
235     def _finalize(group: List[Tuple[int, float, set]]) -> Dict[str, Any]:
236         velocities = [v for _, v, _ in group]
237         track_ids = set().union(*[ids for _, _, ids in group])
238         return {
239             "start": group[0][0] / fps + chunk_start,
240             "end": group[-1][0] / fps + chunk_start,
241             "active_frame_count": len(group),
242             "peak_velocity": max(velocities),
243             "mean_velocity": sum(velocities) / len(velocities),
244             "track_id_count": len(track_ids),
245             "chunk_index": chunk_index,
246         }
247
248     windows = []
249     group = [active_frames[0]]
250     for af in active_frames[1:]:
251         if af[0] - group[-1][0] <= max_gap_frames:
252             group.append(af)
253         else:
254             windows.append(_finalize(group))
255             group = [af]
256     windows.append(_finalize(group))
257
258     # Filter out extremely short action windows
259     min_window_duration = self.config.get("indexing",
{ }).get("min_action_window_duration", 2.0)
260     windows = [w for w in windows if (w["end"] - w["start"]) >= min_window_duration]
261
262     return windows

```

```

263         def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None, max_frames: int = None) -> List[Dict[str, Any]]:
264             self._lazy_load_model()
265
266             # Get sport profile config
267             sport_name = self.config.get("sport", "badminton")
268             try:
269                 sport_profile = sport_registry.get(sport_name)
270             except KeyError as e:
271                 logger.error(str(e))
272                 return []
273
274             # Get AI provider and model config
275             idx_cfg = self.config.get("indexing", {})
276             provider_name = self.config.get("ai_provider", idx_cfg.get("ai_provider",
"gemini"))
277             model_name = self.config.get("ai_model", idx_cfg.get("ai_model",
"gemini-2.5-flash"))
278
279             # Get appropriate API key based on the provider
280             api_key_env_var = f"{provider_name.upper()}_API_KEY"
281             api_key = os.environ.get(api_key_env_var)
282             if not api_key:
283                 logger.error(
284                     f"{api_key_env_var} environment variable is not set! "
285                     f"To run the {provider_name} segmenter, set your API key. "
286                     f"In PowerShell: $env:{api_key_env_var}=\"your_api_key_here\""
287                 )
288                 return []
289
290             try:
291                 provider = provider_registry.get(provider_name, api_key=api_key,
model_name=model_name)
292             except KeyError as e:
293                 logger.error(str(e))
294                 return []
295
296             logger.info(f"Initializing Hybrid Segmenter (torchvision + ByteTrack) for video:
{video_path}")
297
298             video_duration = get_video_duration(video_path)
299             if video_duration <= 0:
300                 logger.error("Invalid video duration.")
301                 return []
302
303             chunk_duration = idx_cfg.get("chunk_duration", 300)
304             chunk_overlap = idx_cfg.get("chunk_overlap", 30)
305             velocity_thresh = idx_cfg.get("yolo_velocity_threshold", 0.15) # Fraction of
frame width per second
306             merge_gap = idx_cfg.get("dead_time_merge_gap", 2.0)
307             frame_skip = idx_cfg.get("yolo_frame_skip", 3)
308             tracker_config = idx_cfg.get("yolo_tracker", "bytetrack.yaml")
309             handoff_padding = idx_cfg.get("ai_handoff_padding",
idx_cfg.get("gemini_handoff_padding", 2.0))
310             ai_max_workers = idx_cfg.get("ai_max_workers", idx_cfg.get("gemini_max_workers",
5))
311             log_candidates = idx_cfg.get("log_yolo_candidates", True)
312             store_candidates = idx_cfg.get("store_yolo_candidates", True)
313
314             # Snapshot of the params that produced these detections ? stored alongside each
315             # candidate so a fine-tuning run can reproduce / correlate against them.
316             detection_params = {
317                 "velocity_thresh": velocity_thresh,
318                 "merge_gap": merge_gap,
319                 "frame_skip": frame_skip,

```



```

320         "tracker": tracker_config,
321         "detector_model": idx_cfg.get("detector_model", _DEFAULT_DETECTOR),
322         "detector_score_threshold": idx_cfg.get("detector_score_threshold", 0.5),
323         "min_action_window_duration": idx_cfg.get("min_action_window_duration",
324 2.0),
325     }
326     chunks_dir = os.path.join("output", "chunks_yolo")
327     os.makedirs(chunks_dir, exist_ok=True)
328
329     step = chunk_duration - chunk_overlap
330     chunk_starts = []
331     t = 0.0
332     while t < video_duration:
333         chunk_starts.append(t)
334         t += step
335
336     num_chunks = len(chunk_starts)
337     logger.info(f"Video duration: {video_duration:.1f}s. Splitting into {num_chunks}
chunks.")
338
339     all_candidate_windows = []
340
341     t_start = time.time()
342     prefetch_workers = idx_cfg.get("yolo_prefetch_workers", 2)
343
344     if num_chunks > 1 and prefetch_workers > 0:
345         # --- PIPELINED PROCESSING ---
346         # GPU inference must stay sequential (single GPU), but ffmpeg chunk
347         # extraction is pure I/O. We pre-extract upcoming chunks in background
348         # threads so extraction of chunk N+1 overlaps with inference on chunk N.
349         logger.info(f"Running detection Phase 2 with pipelined extraction
({prefetch_workers} prefetch workers)...")
350
351         extract_executor =
concurrent.futures.ThreadPoolExecutor(max_workers=prefetch_workers)
352
353         def _extract_chunk(ci: int, cs: float) -> Tuple[int, float, str, bool]:
354             chunk_end = min(cs + chunk_duration, video_duration)
355             actual_dur = chunk_end - cs
356             chunk_path = os.path.join(chunks_dir, f"chunk_{video_id}_{ci}.mp4")
357             success = extract_video_chunk(video_path, cs, actual_dur, chunk_path)
358             return (ci, cs, chunk_path, success)
359
360         # Submit all extractions upfront ? they'll queue and run as workers free up
361         extract_futures = [
362             extract_executor.submit(_extract_chunk, ci, cs)
363             for ci, cs in enumerate(chunk_starts)
364         ]
365
366         # Process results in order: wait for extraction, then run inference on GPU
367         for future in extract_futures:
368             ci, chunk_start, chunk_path, success = future.result()
369             chunk_end = min(chunk_start + chunk_duration, video_duration)
370             logger.info(f"Processing Chunk {ci+1}/{num_chunks} ({chunk_start:.1f}s -
{chunk_end:.1f}s)...")
371
372             if not success:
373                 logger.warning(f"Chunk {ci+1} extraction failed. Skipping.")
374                 continue
375
376             windows = self._extract_action_windows_from_chunk(
377                 chunk_path, chunk_start, velocity_thresh, merge_gap,
378                 frame_skip=frame_skip, tracker_config=tracker_config, chunk_index=ci
379             )

```

```

380         logger.info(f"Found {len(windows)} candidate action windows in chunk
{ci+1}.")
381         all_candidate_windows.extend(windows)
382
383         try:
384             os.remove(chunk_path)
385         except Exception:
386             pass
387
388         extract_executor.shutdown(wait=False)
389     else:
390         # --- SEQUENTIAL PROCESSING (single chunk or prefetch disabled) ---
391         for ci, chunk_start in enumerate(chunk_starts):
392             chunk_end = min(chunk_start + chunk_duration, video_duration)
393             actual_dur = chunk_end - chunk_start
394             chunk_path = os.path.join(chunks_dir, f"chunk_{video_id}_{ci}.mp4")
395
396             logger.info(f"Processing Chunk {ci+1}/{num_chunks} ({chunk_start:.1f}s -
{chunk_end:.1f}s)...")
397             success = extract_video_chunk(video_path, chunk_start, actual_dur,
chunk_path)
398             if not success:
399                 continue
400
401             windows = self._extract_action_windows_from_chunk(
402                 chunk_path, chunk_start, velocity_thresh, merge_gap,
403                 frame_skip=frame_skip, tracker_config=tracker_config, chunk_index=ci
404             )
405             logger.info(f"Found {len(windows)} candidate action windows in chunk
{ci+1}.")
406             all_candidate_windows.extend(windows)
407
408             try:
409                 os.remove(chunk_path)
410             except Exception:
411                 pass
412
413         # Deduplicate overlapping candidate windows across chunks (chunks overlap by
414         # chunk_overlap, so the same rally can surface in two adjacent chunks).
415         def _merge_stats(a: Dict[str, Any], b: Dict[str, Any]) -> Dict[str, Any]:
416             fa, fb = a["active_frame_count"], b["active_frame_count"]
417             total = fa + fb or 1
418             return {
419                 "start": a["start"],
420                 "end": max(a["end"], b["end"]),
421                 "active_frame_count": fa + fb,
422                 "peak_velocity": max(a["peak_velocity"], b["peak_velocity"]),
423                 # Frame-weighted mean so longer sub-windows dominate proportionally.
424                 "mean_velocity": (a["mean_velocity"] * fa + b["mean_velocity"] * fb) /
total,
425                 "track_id_count": max(a["track_id_count"], b["track_id_count"]),
426                 "chunk_index": a["chunk_index"],
427             }
428
429         if all_candidate_windows:
430             all_candidate_windows.sort(key=lambda w: w["start"])
431             merged_windows = [all_candidate_windows[0]]
432             for current in all_candidate_windows[1:]:
433                 last = merged_windows[-1]
434                 if current["start"] <= last["end"]: # Overlap
435                     merged_windows[-1] = _merge_stats(last, current)
436                 else:
437                     merged_windows.append(current)
438             all_candidate_windows = merged_windows
439

```

```

440         logger.info(f"Phase 2 tracking completed in {time.time() - t_start:.1f}s.")
441         logger.info(f"Total candidate windows: {len(all_candidate_windows)}")
442
443     # Per-rally console log (final, full-video timestamps) for video cross-
verification.
444     if log_candidates:
445         for w in all_candidate_windows:
446             logger.info(
447                 f"[rally] {_fmt_ts(w['start'])}-{_fmt_ts(w['end'])} "
448                 f"({w['end'] - w['start']:.1f}s) | frames={w['active_frame_count']} "
449                 f"peak_v={w['peak_velocity']:.3f} mean_v={w['mean_velocity']:.3f} "
450                 f"tracks={w['track_id_count']}"
451             )
452
453     # Persist raw candidates for the fine-tuning dataset. Map window index ->
candidate_id
454     # so confirmed AI segments can be linked back in Phase 4.
455     candidate_ids = [None] * len(all_candidate_windows)
456     if store_candidates:
457         for i, w in enumerate(all_candidate_windows):
458             stats = {
459                 "peak_velocity": w["peak_velocity"],
460                 "mean_velocity": w["mean_velocity"],
461                 "active_frame_count": w["active_frame_count"],
462                 "track_id_count": w["track_id_count"],
463             }
464             # Use peak velocity (capped at 1.0) as a coarse generic confidence
score.
465             confidence = min(1.0, w["peak_velocity"])
466             candidate_ids[i] = self.db.add_candidate_segment(
467                 video_id, detector="yolo_hybrid",
468                 start_time=w["start"], end_time=w["end"],
469                 chunk_index=w["chunk_index"], confidence=confidence,
470                 stats=stats, detection_params=detection_params,
471             )
472
473     try:
474         os.rmdir(chunks_dir)
475     except Exception:
476         pass
477
478     # Phase 3: AI Provider Handoff
479     gemini_segments = []
480     gemini_dir = os.path.join("output", f"chunks_{provider_name}_handoff")
481     os.makedirs(gemini_dir, exist_ok=True)
482
483     logger.info(f"Starting Phase 3: AI Provider ({provider_name}) Validation on
{len(all_candidate_windows)} candidate windows...")
484
485     def process_candidate_window(wi: int, window: Dict[str, Any]) -> List[dict]:
486         w_start = max(0, window["start"] - handoff_padding)
487         w_end = min(video_duration, window["end"] + handoff_padding)
488         w_dur = w_end - w_start
489
490         w_path = os.path.join(gemini_dir, f"handoff_{video_id}_{wi}.mp4")
491         # Re-encode short handoff clips to avoid keyframe-boundary drift
492         success = extract_video_chunk(video_path, w_start, w_dur, w_path,
reencode=True)
493         if not success:
494             return []
495
496         prompt = sport_profile.get_prompt(chunk_duration=w_dur)
497         segments, metrics = provider.analyze_video(w_path, prompt,
chunk_duration=w_dur, label=f"Handoff {wi+1}")

```

```

498
499     # Coordinate mapping ? shift local timestamps back to full-video time
500     valid_segments = []
501     for segment in segments:
502         try:
503             segment["start_time"] = float(segment["start_time"]) + w_start
504             segment["end_time"] = float(segment["end_time"]) + w_start
505             # Stamp the source candidate so Phase 4 can link
506             prediction->confirmation.
507             segment["_candidate_id"] = candidate_ids[wi]
508             valid_segments.append(segment)
509             except (ValueError, TypeError, KeyError) as e:
510                 logger.warning(f"Dropping segment with unparseable timestamps:
511 {segment} ? {e}")
512
513         try:
514             os.remove(w_path)
515         except OSError:
516             pass
517
518     return valid_segments
519
520     with concurrent.futures.ThreadPoolExecutor(max_workers=ai_max_workers) as
521 executor:
522     futures = [executor.submit(process_candidate_window, wi, w) for wi, w in
523 enumerate(all_candidate_windows)]
524     # Iterate in submission order to preserve temporal ordering
525     for future in futures:
526         gemini_segments.extend(future.result())
527
528     try:
529         os.rmdir(gemini_dir)
530     except OSError:
531         pass
532
533     logger.info(f"Phase 3 completed. AI confirmed {len(gemini_segments)} semantic
534 play segments.")
535
536     # Phase 4: DB Population
537     segments_data = []
538     for segment in gemini_segments:
539         try:
540             start = float(segment["start_time"])
541             end = float(segment["end_time"])
542             dur = end - start
543
544             shot_count = segment.get("shuttle_exchange_count")
545             if shot_count is None: shot_count = max(3, int(dur / 0.8))
546             else:
547                 try: shot_count = int(shot_count)
548                 except (ValueError, TypeError): shot_count = max(3, int(dur / 0.8))
549
550             metadata = {
551                 "shot_count": shot_count,
552                 "shot_count_confidence": segment.get("shot_count_confidence",
553 "Unknown"),
554                 "detector": f"yolo-hybrid-{model_name}"
555             }
556
557             segment_id = self.db.add_play_segment(video_id, start, end, None, None,
558 metadata)
559
560             if segment_id:
561                 # Link back to the candidate that produced this segment. A candidate
562                 # may yield multiple AI segments (one-to-many); confirmed_segment_id
563                 keeps

```

```

555             # the last one ? the link table records "this candidate was
confirmed".
556             candidate_id = segment.get("_candidate_id")
557             if candidate_id is not None:
558                 self.db.link_candidate_to_segment(candidate_id, segment_id)
559             segments_data.append({
560                 "id": segment_id,
561                 "start_time": start,
562                 "end_time": end,
563                 "duration": dur,
564                 "metadata": metadata
565             })
566         except Exception as e:
567             logger.warning(f"Skipping invalid segment struct: {e}")
568
569     return segments_data
570
571 # Register the segmenter
572 segmenter_registry.register(HybridYoloSegmenter)
573

```

18. user (2026-06-10T06:02:21.270Z)

```

1     import os
2     import time
3     import json
4     import logging
5     import subprocess
6     import concurrent.futures
7     from typing import List, Dict, Any, Tuple
8     from backend.infrastructure.database import Database
9     from backend.pipeline.segmenters.base import BasePlaySegmenter, segmenter_registry
10    from backend.utils.video import get_video_duration, extract_video_chunk
11    from backend.sports import sport_registry
12    from backend.providers import provider_registry
13
14    logger = logging.getLogger(__name__)
15
16    class GeminiPlaySegmenter(BasePlaySegmenter):
17        @property
18        def name(self) -> str:
19            return "gemini"
20
21        @property
22        def description(self) -> str:
23            return "High-precision AI-powered semantic play segment detection using Google
Gemini Multimodal API."
24
25        def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None, max_frames: int = None) -> List[Dict[str, Any]]:
26            """
27            Processes a video to detect play segments using Google Gemini Multimodal Video
API,
28            populates SQLite, and indexes detailed events.
29            """
30            # Get sport profile config
31            sport_name = self.config.get("sport", "badminton")
32            try:
33                sport_profile = sport_registry.get(sport_name)
34            except KeyError as e:
35                logger.error(str(e))
36            return []
37
38            # Get AI provider and model config
39            idx_cfg = self.config.get("indexing", {})

```

```

40         provider_name = self.config.get("ai_provider", idx_cfg.get("ai_provider",
"gemini"))
41         model_name = self.config.get("ai_model", idx_cfg.get("ai_model",
"gemini-2.5-flash"))
42
43         # Get appropriate API key based on the provider
44         api_key_env_var = f"{provider_name.upper()}_API_KEY"
45         api_key = os.environ.get(api_key_env_var)
46         if not api_key:
47             logger.error(
48                 f"{api_key_env_var} environment variable is not set! "
49                 f"To run the {provider_name} segmenter, set your API key. "
50                 f"In PowerShell: $env:{api_key_env_var}=\"your_api_key_here\""
51             )
52             return []
53
54         try:
55             provider = provider_registry.get(provider_name, api_key=api_key,
model_name=model_name)
56         except KeyError as e:
57             logger.error(str(e))
58             return []
59
60         # A: Check for debug_duration to trim the video first
61         debug_duration = self.config.get("indexing", {}).get("debug_duration")
62         video_to_upload = video_path
63         is_trimmed = False
64         temp_trimmed_path = None
65         is_compressed = False
66         temp_compressed_path = None
67
68         if debug_duration:
69             logger.info(f"DEBUG FEATURE: Trimming video to first {debug_duration}
seconds for fast processing...")
70             temp_trimmed_path = os.path.join("output", f"temp_trimmed_{video_id}.mp4")
71             os.makedirs("output", exist_ok=True)
72             if os.path.exists(temp_trimmed_path):
73                 try:
74                     os.remove(temp_trimmed_path)
75                 except Exception:
76                     pass
77
78             cmd = [
79                 "ffmpeg", "-y",
80                 "-ss", "0",
81                 "-t", str(debug_duration),
82                 "-i", video_path,
83                 "-c", "copy",
84                 temp_trimmed_path
85             ]
86             logger.debug(f"Running command: {' '.join(cmd)}")
87             try:
88                 subprocess.run(cmd, check=True, stdout=subprocess.DEVNULL,
stderr=subprocess.DEVNULL)
89                 logger.info(f"Trimmed video successfully created at:
{temp_trimmed_path}")
90                 video_to_upload = temp_trimmed_path
91                 is_trimmed = True
92             except Exception as e:
93                 logger.warning(f"Failed to trim video using ffmpeg: {e}. Falling back to
full video.")
94
95
96
97         # Helper to clean up local temp files

```

```

98     def cleanup_temp_files():
99         if is_trimmed and temp_trimmed_path and os.path.exists(temp_trimmed_path):
100             try:
101                 os.remove(temp_trimmed_path)
102             except Exception:
103                 pass
104
105     # 1. Determine actual video duration
106     video_duration = get_video_duration(video_to_upload)
107     if video_duration > 0:
108         logger.info(f"Detected video duration: {video_duration:.1f}s")
109     else:
110         logger.warning("Could not detect video duration. Prompt will not include
duration context.")
111
112     # 2. Decide whether to use chunked processing
113     idx_cfg = self.config.get("indexing", {})
114     chunk_duration = idx_cfg.get("chunk_duration", 300)    # 5 minutes default
115     chunk_overlap = idx_cfg.get("chunk_overlap", 30)      # 30 seconds overlap
116     use_chunking = video_duration > 0 and video_duration > chunk_duration
117
118     ai_segments = []
119     chunks_dir = os.path.join("output", "chunks")
120
121     if use_chunking:
122         # --- CHUNKED PROCESSING ---
123         step = chunk_duration - chunk_overlap
124         chunk_starts = []
125         t = 0.0
126         while t < video_duration:
127             chunk_starts.append(t)
128             t += step
129         num_chunks = len(chunk_starts)
130         logger.info(
131             f"Video is {video_duration:.1f}s ? splitting into {num_chunks}
overlapping chunks "
132             f"({chunk_duration}s each, {chunk_overlap}s overlap).")
133         )
134
135         os.makedirs(chunks_dir, exist_ok=True)
136
137         def process_single_chunk(ci: int, chunk_start: float) -> Tuple[str, list,
dict]:
138             chunk_end = min(chunk_start + chunk_duration, video_duration)
139             actual_chunk_dur = chunk_end - chunk_start
140             chunk_label = f"Chunk {ci+1}/{num_chunks}"
141             chunk_path = os.path.join(chunks_dir, f"chunk_{video_id}_{ci}.mp4")
142
143             # Extract chunk with ffmpeg
144             logger.info(f"Extracting {chunk_label}: {chunk_start:.1f}s -
{chunk_end:.1f}s ({actual_chunk_dur:.1f}s)")
145             success = extract_video_chunk(video_to_upload, chunk_start,
actual_chunk_dur, chunk_path)
146             if not success:
147                 logger.error(f"Failed to extract {chunk_label}. Skipping.")
148                 return chunk_label, [], {}
149
150             # Measure chunk size
151             try:
152                 chunk_size_mb = os.path.getsize(chunk_path) / (1024 * 1024)
153             except Exception:
154                 chunk_size_mb = 0.0
155
156             # Get prompt from sport profile
157             prompt = sport_profile.get_prompt(chunk_duration=actual_chunk_dur)

```

```

158
159         # Process this chunk through the provider
160         chunk_segments, metrics = provider.analyze_video(
161             chunk_path, prompt, chunk_duration=actual_chunk_dur,
label=chunk_label
162         )
163         metrics["file_size_mb"] = chunk_size_mb
164         metrics["segments_found"] = len(chunk_segments)
165
166         # Offset timestamps back to full-video coordinates
167         for segment in chunk_segments:
168             if "start_time" in segment:
169                 try:
170                     segment["start_time"] = float(segment["start_time"]) +
chunk_start
171                 except (ValueError, TypeError):
172                     pass
173             if "end_time" in segment:
174                 try:
175                     segment["end_time"] = float(segment["end_time"]) +
chunk_start
176                 except (ValueError, TypeError):
177                     pass
178
179         logger.info(f"{chunk_label} returned {len(chunk_segments)} play
segments.")
180
181         # Clean up chunk file
182         try:
183             os.remove(chunk_path)
184         except Exception:
185             pass
186
187         return chunk_label, chunk_segments, metrics
188
189     # Run chunks in parallel
190     all_metrics = {}
191     logger.info("Starting parallel processing with up to 5 concurrent
workers...")
192     t_chunking_start = time.time()
193     with concurrent.futures.ThreadPoolExecutor(max_workers=5) as executor:
194         futures = [executor.submit(process_single_chunk, ci, start) for ci,
start in enumerate(chunk_starts)]
195         for future in concurrent.futures.as_completed(futures):
196             label, segments, metrics = future.result()
197             ai_segments.extend(segments)
198             if metrics:
199                 all_metrics[label] = metrics
200
201     # Print Summary Log
202     summary_lines = [
203         "",
204         "-" * 60,
205         "                                GEMINI CHUNKING SUMMARY",
206         "-" * 60,
207         f" Total time: {time.time() - t_chunking_start:.1f}s",
208         f" Total chunks: {num_chunks}",
209         f" Total play segments found: {len(ai_segments)}",
210         "-" * 60,
211         f" {'Chunk':<12} | {'Size':<8} | {'Upload':<8} | {'Process':<8} |
{'Gen.':<8} | {'Segments':<8}",
212         "-" * 60,
213     ]
214     for i in range(num_chunks):
215         lbl = f"Chunk {i+1}/{num_chunks}"

```



```

216         m = all_metrics.get(lbl, {})
217         sz = f"{m.get('file_size_mb', 0):.1f}MB" if m else "N/A"
218         upl = f"{m.get('upload_time', 0):.1f}s" if m else "N/A"
219         prc = f"{m.get('process_time', 0):.1f}s" if m else "N/A"
220         gen = f"{m.get('gen_time', 0):.1f}s" if m else "N/A"
221         seg = m.get('segments_found', 0) if m else 0
222         summary_lines.append(f" {lbl:<12} | {sz:<8} | {upl:<8} | {prc:<8} |
{gen:<8} | {seg}>")
223     summary_lines.append("=" * 60)
224     logger.info("\n".join(summary_lines))
225
226     # Clean up local temp files
227     cleanup_temp_files()
228     try:
229         os.rmdir(chunks_dir)
230     except Exception:
231         pass
232
233     logger.info(f"All chunks processed. {len(ai_segments)} total play segments
before deduplication.")
234
235     else:
236         # --- SINGLE-FILE PROCESSING (short videos) ---
237         try:
238             file_sz = os.path.getsize(video_to_upload) / (1024 * 1024)
239         except Exception:
240             file_sz = 0.0
241
242         # Get prompt from sport profile
243         prompt = sport_profile.get_prompt(chunk_duration=video_duration)
244         ai_segments, metrics = provider.analyze_video(
245             video_to_upload, prompt, chunk_duration=video_duration,
label="SingleFile"
246         )
247         summary_lines = [
248             "",
249             "=" * 50,
250             "      GEMINI SINGLE FILE SUMMARY",
251             "=" * 50,
252             f" File Size: {file_sz:.1f} MB",
253             f" Upload time: {metrics.get('upload_time', 0):.1f}s",
254             f" API Processing time: {metrics.get('process_time', 0):.1f}s",
255             f" Prompt Generation time: {metrics.get('gen_time', 0):.1f}s",
256             f" Play segments found: {len(ai_segments)}",
257             "=" * 50,
258         ]
259         logger.info("\n".join(summary_lines))
260         cleanup_temp_files()
261
262     if not ai_segments:
263         return []
264
265     # Populate SQLite DB with parsed play segments
266     segments_data = []
267     logger.info(f"Processing {len(ai_segments)} play segments through
validation...")
268
269     def parse_time(val) -> float:
270         """Converts a timestamp value to total float seconds.
271         Handles float/int pass-through, plain numeric strings, and
272         colon-separated formats (MM:SS, HH:MM:SS). Logs warnings
273         when unexpected formats are encountered."""
274         if isinstance(val, (int, float)):
275             return float(val)
276         if isinstance(val, str):

```

```

277         val = val.strip()
278         if ":" in val:
279             # Gemini returned MM:SS or HH:MM:SS despite being asked for decimal
seconds.
280             # Convert it, but warn so the issue is visible in logs.
281             parts = val.split(":")
282             parts.reverse()
283             total = 0.0
284             for i, p in enumerate(parts):
285                 try:
286                     total += float(p) * (60 ** i)
287                 except ValueError:
288                     pass
289                 logger.warning(f"Timestamp '{val}' was in colon-separated format
(MM:SS/HH:MM:SS) instead of decimal seconds. Converted to {total:.2f}s.")
290             return total
291         try:
292             return float(val)
293         except ValueError:
294             logger.warning(f"Could not parse timestamp string '{val}' as a
number. Defaulting to 0.0.")
295             return 0.0
296         logger.warning(f"Unexpected timestamp type {type(val).__name__} for value
'{val}'. Defaulting to 0.0.")
297         return 0.0
298
299     # 7. Parse timestamps and build candidate segment list
300     idx_cfg = self.config.get("indexing", {})
301     min_segment_duration = idx_cfg.get("min_segment_duration", 3.0)
302     max_segment_duration = 90.0 # Singles play segments almost never exceed this
303
304     candidates = []
305     for idx, segment in enumerate(ai_segments):
306         start = parse_time(segment.get("start_time", 0.0))
307         end = parse_time(segment.get("end_time", 0.0))
308         if start >= end:
309             logger.info(f"Skipping segment #{idx+1}: start_time ({start:.2f}s) >=
end_time ({end:.2f}s)")
310             continue
311             candidates.append({
312                 "idx": idx,
313                 "start": start,
314                 "end": end,
315                 "raw": segment
316             })
317
318     # 8. Post-hoc validation pass
319     logger.info(f"Running validation on {len(candidates)} candidate play
segments...")
320     validated = []
321     rejected = []
322
323     for c in candidates:
324         label = f"Segment #{c['idx']+1} [{c['start']:.2f}s - {c['end']:.2f}s]"
325
326         # A: Reject segments that start beyond the video
327         if video_duration > 0 and c["start"] >= video_duration:
328             rejected.append((label, f"start_time ({c['start']:.2f}s) is beyond video
duration ({video_duration:.1f}s)"))
329             continue
330
331         # B: Clamp end_time to video duration
332         if video_duration > 0 and c["end"] > video_duration:
333             logger.info(f"{label}: end_time ({c['end']:.2f}s) exceeds video duration
({video_duration:.1f}s). Clamping.")

```

```

334         c["end"] = video_duration
335
336         dur = c["end"] - c["start"]
337
338         # C: Reject segments shorter than min_segment_duration
339         if dur < min_segment_duration:
340             rejected.append((label, f"duration ({dur:.2f}s) is below minimum
({min_segment_duration}s)"))
341             continue
342
343         # D: Flag suspiciously long segments
344         if dur > max_segment_duration:
345             logger.warning(
346                 f"{label}: duration ({dur:.1f}s) is unusually long
(>{max_segment_duration}s). "
347                 f"This may be a detection error. Keeping but flagging."
348             )
349
350         validated.append(c)
351
352     # E: Resolve overlapping segments (keep the longer one)
353     if len(validated) > 1:
354         # Sort by start time
355         validated.sort(key=lambda r: r["start"])
356         non_overlapping = [validated[0]]
357         for curr in validated[1:]:
358             prev = non_overlapping[-1]
359             # Check if current segment overlaps with the previous one
360             if curr["start"] < prev["end"]:
361                 prev_dur = prev["end"] - prev["start"]
362                 curr_dur = curr["end"] - curr["start"]
363                 kept, dropped = (prev, curr) if prev_dur >= curr_dur else (curr,
prev)
364                 kept_label = f"Segment #{kept['idx']+1} [{kept['start']:.2f}s -
{kept['end']:.2f}s]"
365                 dropped_label = f"Segment #{dropped['idx']+1}
[{dropped['start']:.2f}s - {dropped['end']:.2f}s]"
366                 rejected.append((dropped_label, f"overlaps with {kept_label} ? kept
the longer segment"))
367                 if kept == curr:
368                     non_overlapping[-1] = curr # Replace previous with current
369                 else:
370                     non_overlapping.append(curr)
371             validated = non_overlapping
372
373     # Log validation summary
374     if rejected:
375         reject_lines = [f" ? {label}: {reason}" for label, reason in rejected]
376         logger.info(f"{len(rejected)} play segments were REJECTED:\n" +
"\n".join(reject_lines))
377         logger.info(f"{len(validated)} play segments passed validation (out of
{len(ai_segments)} from Gemini).")
378
379     stitch_cfg = self.config.get("stitching", {})
380     begin_pad = stitch_cfg.get("begin_padding", 0.5)
381     end_pad = stitch_cfg.get("end_padding", 0.5)
382
383     total_highlight_dur = 0.0
384     for c in validated:
385         total_highlight_dur += (c["end"] - c["start"]) + begin_pad + end_pad
386
387     mins = int(total_highlight_dur // 60)
388     secs = int(total_highlight_dur % 60)
389     logger.info(
390         f"Estimated final highlight video duration: {mins:02d}:{secs:02d} "

```

```

391         f"({total_highlight_dur:.1f}s total; includes {begin_pad}s begin +
{end_pad}s end padding per segment)"
392     )
393
394     # 9. Populate SQLite DB with validated play segments
395     for c in validated:
396         start = c["start"]
397         end = c["end"]
398         segment = c["raw"]
399         dur = end - start
400
401         # Extract shot_count from Gemini or compute estimate if not present
402         shot_count = segment.get("shuttle_exchange_count")
403         if shot_count is None:
404             shot_count = max(3, int(dur / 0.8)) # 1 shot every 0.8 seconds as
fallback estimate
405         else:
406             try:
407                 shot_count = int(shot_count)
408             except Exception:
409                 shot_count = max(3, int(dur / 0.8))
410
411         # Only store metrics we actually know ? no fabricated data
412         metadata = {
413             "shot_count": shot_count,
414             "shot_count_confidence": segment.get("shot_count_confidence",
"Unknown"),
415             "detector": model_name
416         }
417
418         segment_id = self.db.add_play_segment(video_id, start, end, None, None,
metadata)
419
420         if segment_id:
421             segments_data.append({
422                 "id": segment_id,
423                 "start_time": start,
424                 "end_time": end,
425                 "duration": dur,
426                 "metadata": metadata
427             })
428
429     return segments_data
430
431 # Register the segmenter
432 segmenter_registry.register(GeminiPlaySegmenter)
433

```

19. user (2026-06-10T06:02:43.673Z)

```

1      {
2          "sport": "badminton",
3          "ai_provider": "gemini",
4          "ai_model": "gemini-2.5-flash",
5          "validation": {
6              "min_resolution_width": 1280,
7              "min_resolution_height": 720,
8              "min_fps": 29.9,
9              "min_blur_threshold": 20.0,
10             "max_scene_cuts_per_minute": 0.5,
11             "relevance": {
12                 "court_green_lower": [
13                     35,
14                     40,
15                     40

```

```

16         ],
17         "court_green_upper": [
18             85,
19             255,
20             255
21         ],
22         "court_pixels_min_ratio": 0.05
23     }
24 },
25     "indexing": {
26         "default_segmenter": "gemini",
27         "use_gpu": true,
28         "motion_threshold": 0.08,
29         "min_segment_duration": 3.0,
30         "dead_time_merge_gap": 2.0,
31         "chunk_duration": 90,
32         "chunk_overlap": 10,
33         "detector_model": "fasterrcnn_resnet50_fpn",
34         "detector_score_threshold": 0.5,
35         "yolo_velocity_threshold": 0.15,
36         "yolo_frame_skip": 3,
37         "yolo_tracker": "bytetrack.yaml",
38         "yolo_parallel_chunks": false,
39         "yolo_chunk_workers": 2,
40         "min_action_window_duration": 2.0,
41         "log_yolo_candidates": true,
42         "store_yolo_candidates": true,
43         "ai_handoff_padding": 2.0,
44         "ai_max_workers": 5,
45         "tracknet": {
46             "wsl_distro": "Ubuntu",
47             "repo_dir": "~/models/TrackNetV4",
48             "conda_sh": "~/miniconda3/etc/profile.d/conda.sh",
49             "conda_env": "TrackNetV4",
50             "weights_path": "",
51             "queue_length": 5,
52             "stage_in_wsl": true,
53             "wsl_stage_dir": "~/clips",
54             "timeout_sec": 0,
55             "velocity_threshold": 0.02
56         }
57     },
58     "output": {
59         "default_highlights_name": "highlights.mp4",
60         "db_name": "sports_indexer.db"
61     },
62     "stitching": {
63         "begin_padding": 0.5,
64         "end_padding": 0.5,
65         "use_cache": false,
66         "min_shot_count": 1
67     }
68 }

```

20. user (2026-06-10T06:02:44.260Z)

```

1     """WASB shuttle-trajectory runner (WSL2 subprocess adapter).
2
3     Windows-side counterpart to backend/detectors/wasb_infer.py. It:
4     1. stages the video into the WSL filesystem (fast local reads, not /mnt),
5     2. copies wasb_infer.py into the WASB repo's src/ dir,
6     3. runs it in the WSL `wasb` conda env -> a Frame,Visibility,X,Y trajectory CSV,
7     4. reuses the (model-agnostic) trajectory->action-window logic from tracknet_runner.
8
9     So the only WASB-specific bits live here; the windowing + CSV parsing are shared

```

```

10     with the TrackNet path. Verified inference core: wasb_infer.py (see its docstring).
11     """
12     import os
13     import shlex
14     import logging
15     import subprocess
16     from dataclasses import dataclass
17     from typing import Dict, Any, Optional, Tuple
18
19     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical.
20     # parse_trajectory_csv / trajectory_to_action_windows are @staticmethods on
21     TrackNetRunner.
22     from backend.pipeline.detectors.tracknet_runner import to_wsl_mnt_path, TrackNetRunner
23     from backend.pipeline.detectors.base import DetectorRunner
24
25     logger = logging.getLogger(__name__)
26
27     # Windows path to the inference wrapper we stage into WSL.
28     _INFER_WIN = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
29
30     @dataclass
31     class WasbConfig:
32         """Resolved settings for a WASB WSL run (built from config.json ->
33         indexing.wasb)."""
34         distro: str = "Ubuntu"
35         repo_dir: str = "~/models/WASB-SBDT"
36         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
37         conda_env: str = "wasb"
38         weights_path: str = "~/models/WASB-
39         SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
40         sport: str = "badminton"
41         wsl_stage_dir: str = "~/clips"
42         timeout_sec: int = 0 # 0 = no timeout
43
44     @classmethod
45     def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "WasbConfig":
46         w = (idx_cfg or {}).get("wasb", {}) or {}
47         return cls(
48             distro=w.get("wsl_distro", "Ubuntu"),
49             repo_dir=w.get("repo_dir", "~/models/WASB-SBDT"),
50             conda_sh=w.get("conda_sh", "~/miniconda3/etc/profile.d/conda.sh"),
51             conda_env=w.get("conda_env", "wasb"),
52             weights_path=w.get("weights_path",
53                               "~/models/WASB-
54                               SBDT/pretrained_weights/wasb_badminton_best.pth.tar"),
55             sport=w.get("sport", "badminton"),
56             wsl_stage_dir=w.get("wsl_stage_dir", "~/clips"),
57             timeout_sec=int(w.get("timeout_sec", 0)),
58         )
59
60     class WasbRunner(DetectorRunner): # DetectorRunner ABC (finish item 6 / new item 1)
61         def __init__(self, cfg: WasbConfig):
62             self.cfg = cfg
63
64         def _wsl_bash(self, inner_cmd: str) -> subprocess.CompletedProcess:
65             full = f"source {self.cfg.conda_sh} && {inner_cmd}"
66             argv = ["wsl", "-d", self.cfg.distro, "bash", "-c", full]
67             logger.debug("WSL exec: %s", full)
68             return subprocess.run(argv, capture_output=True, text=True,
69                                  timeout=(self.cfg.timeout_sec or None))
70
71         def healthcheck(self) -> Tuple[bool, str]:
72             """Verify the WSL `wasb` env imports torch and sees a GPU. Returns (ok, msg)."""

```

```

71         cmd = (f"conda activate {self.cfg.conda_env} && "
72                f"python -c \"import torch; print('torch', torch.__version__, "
73                f"'cuda', torch.cuda.is_available())\"")
74         try:
75             res = self._wsl_bash(cmd)
76         except FileNotFoundError:
77             return False, "`wsl` not found ? Windows + WSL2 required."
78         except subprocess.TimeoutExpired:
79             return False, "WSL healthcheck timed out."
80         out = (res.stdout or "").strip()
81         if res.returncode != 0:
82             return False, f"WSL `wasb` env check failed: {(res.stderr or out).strip()}"
83         return True, out
84
85     def _stage_video(self, video_win_path: str, base: str) -> Optional[str]:
86         src_mnt = to_wsl_mnt_path(video_win_path)
87         dest = f"{self.cfg.wsl_stage_dir}/{base}"
88         cmd = f"mkdir -p {self.cfg.wsl_stage_dir} && cp {shlex.quote(src_mnt)} "
89         {shlex.quote(dest)} && echo OK"
90         try:
91             res = self._wsl_bash(cmd)
92         except subprocess.TimeoutExpired:
93             logger.error("Staging video into WSL timed out.")
94             return None
95         if res.returncode != 0 or "OK" not in (res.stdout or ""):
96             logger.error("Failed to stage video into WSL: %s", (res.stderr or
97 res.stdout).strip())
98             return None
99         return dest
100
101     def _stage_infer_script(self) -> bool:
102         """Copy wasb_infer.py into the WASB repo src/ so it can import WASB modules."""
103         src_mnt = to_wsl_mnt_path(_INFER_WIN)
104         cmd = f"cp {shlex.quote(src_mnt)} {self.cfg.repo_dir}/src/wasb_infer.py && echo
105         OK"
106         res = self._wsl_bash(cmd)
107         return res.returncode == 0 and "OK" in (res.stdout or "")
108
109     def run_predict(self, video_win_path: str, output_win_dir: str) -> Optional[str]:
110         """Run WASB on a video. Returns the Windows path to the trajectory CSV, or
111         None."""
112         os.makedirs(output_win_dir, exist_ok=True)
113         # Normalize for cross-platform basename (tests use Windows paths on Linux).
114         video_win_path = str(video_win_path).replace("\\", "/")
115         base = os.path.basename(video_win_path)
116         stem = os.path.splitext(base)[0]
117         wsl_out_dir = to_wsl_mnt_path(output_win_dir)
118         expected_csv_win = os.path.join(output_win_dir, f"{stem}_wasb.csv")
119
120         if not self._stage_infer_script():
121             logger.error("Could not stage wasb_infer.py into WSL.")
122             return None
123         staged_video = self._stage_video(video_win_path, base)
124         if staged_video is None:
125             return None
126         frames_dir = f"{self.cfg.wsl_stage_dir}/{stem}_frames"
127
128         cmd = (
129             f"conda activate {self.cfg.conda_env} && cd {self.cfg.repo_dir}/src && "
130             f"python wasb_infer.py "
131             f"--video {shlex.quote(staged_video)} "
132             f"--frames_out_dir {shlex.quote(frames_dir)} "
133             f"--weights {shlex.quote(self.cfg.weights_path)} "
134             f"--sport {shlex.quote(self.cfg.sport)} "
135             f"--out {shlex.quote(wsl_out_dir + '/' + stem + '_wasb.csv')}")

```

```

132         )
133         logger.info("Running WASB inference on %s ...", base)
134         try:
135             res = self._wsl_bash(cmd)
136         except subprocess.TimeoutExpired:
137             logger.error("WASB inference timed out after %ss.", self.cfg.timeout_sec)
138             return None
139         if res.returncode != 0:
140             logger.error("WASB inference failed:\n%s", (res.stderr or
res.stdout)[-2000:])
141             return None
142         if not os.path.exists(expected_csv_win):
143             logger.error("WASB finished but expected CSV not found at %s",
expected_csv_win)
144             return None
145         logger.info("WASB trajectory CSV: %s", expected_csv_win)
146         return expected_csv_win
147
148     # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
149     parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
150     trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)
151

```

21. user (2026-06-10T06:03:00.277Z)

```

1     """TrackNet shuttle-trajectory runner (WSL2 subprocess adapter).
2
3     This bridges the Windows-side indexer to a TensorFlow TrackNetV4 install living
4     inside WSL2 (ADR-001 in docs/DECISIONS.md). The flow is:
5
6         1. Translate the Windows video path to its WSL-visible form (and optionally
7            stage it into the WSL filesystem to avoid slow /mnt/c per-frame reads).
8         2. Invoke ``src/predict.py`` inside the TrackNetV4 conda env via ``wsl``.
9         3. Read back the per-frame ``Frame,Visibility,X,Y`` CSV it writes.
10        4. Convert that trajectory into high-motion "action windows" ? the same
11           candidate-rally shape HybridYoloSegmenter produces ? so the rest of the
12           pipeline (AI handoff, DB population) is unchanged.
13
14     Nothing here imports TensorFlow; all model code stays in WSL.
15     """
16
17     import os
18     import csv
19     import shlex
20     import logging
21     import subprocess
22     from dataclasses import dataclass
23     from typing import List, Dict, Any, Optional, Tuple
24
25     from backend.pipeline.detectors.base import DetectorRunner
26
27     logger = logging.getLogger(__name__)
28
29
30     @dataclass
31     class TrackNetConfig:
32         """Resolved settings for a TrackNet WSL run (built from config.json ->
indexing.tracknet)."""
33         distro: str = "Ubuntu"
34         repo_dir: str = "~/models/TrackNetV4" # WSL path to the cloned repo
35         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
36         conda_env: str = "TrackNetV4"
37         weights_path: str = "" # WSL path to the .h5/.keras weights

```



```

(REQUIRED)
38     queue_length: int = 5
39     stage_in_wsl: bool = True                    # copy video into WSL fs first
(perf)
40     wsl_stage_dir: str = "~/clips"                # where staged videos/outputs live
in WSL
41     timeout_sec: int = 0                        # 0 = no timeout
42
43     @classmethod
44     def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "TrackNetConfig":
45         tn = (idx_cfg or {}).get("tracknet", {}) or {}
46         return cls(
47             distro=tn.get("wsl_distro", "Ubuntu"),
48             repo_dir=tn.get("repo_dir", "~/models/TrackNetV4"),
49             conda_sh=tn.get("conda_sh", "~/miniconda3/etc/profile.d/conda.sh"),
50             conda_env=tn.get("conda_env", "TrackNetV4"),
51             weights_path=tn.get("weights_path", ""),
52             queue_length=int(tn.get("queue_length", 5)),
53             stage_in_wsl=bool(tn.get("stage_in_wsl", True)),
54             wsl_stage_dir=tn.get("wsl_stage_dir", "~/clips"),
55             timeout_sec=int(tn.get("timeout_sec", 0)),
56         )
57
58
59     def to_wsl_mnt_path(win_path: str) -> str:
60         """Translate a Windows path (C:\\Users\\x) to its WSL /mnt form (/mnt/c/Users/x).
61
62         Already-POSIX paths (starting with / or ~) are returned unchanged so the
63         same helper is safe to call on either kind of path.
64         """
65         p = str(win_path)
66         if p.startswith("/") or p.startswith("~"):
67             return p
68         p = p.replace("\\", "/")
69         if len(p) >= 2 and p[1] == ":":
70             drive = p[0].lower()
71             return f"/mnt/{drive}{p[2:]}"
72         return p
73
74
75     @dataclass
76     class TrajectoryPoint:
77         frame: int
78         visible: bool
79         x: float
80         y: float
81
82
83     class TrackNetRunner(DetectorRunner): # DetectorRunner ABC (finish item 6 / new item 1)
84         """Runs TrackNet predict.py in WSL and turns the output CSV into action windows.
85
86         Implements the common DetectorRunner contract so a future Linux-native or
87         alternative trajectory runner can be substituted without touching the hybrids.
88         """
89
90         def __init__(self, cfg: TrackNetConfig):
91             self.cfg = cfg
92
93         # ----- WSL exec
94
95         def _wsl_bash(self, inner_cmd: str) -> subprocess.CompletedProcess:
96             """Run a bash command string inside the configured WSL distro."""
97             full = f"source {self.cfg.conda_sh} && {inner_cmd}"
98             argv = ["wsl", "-d", self.cfg.distro, "bash", "-c", full]
99             logger.debug("WSL exec: %s", full)

```

```

100         return subprocess.run(
101             argv, capture_output=True, text=True,
102             timeout=(self.cfg.timeout_sec or None),
103         )
104
105     def healthcheck(self) -> Tuple[bool, str]:
106         """Verify the WSL env is usable: conda env exists, TF imports, GPU visible.
107
108         Returns (ok, message). Cheap to call before a long run.
109         """
110         cmd = (
111             f"conda activate {self.cfg.conda_env} && "
112             f"python -c \"import tensorflow as tf; "
113             f"print('TF', tf.__version__, 'GPUs', "
114             f"len(tf.config.list_physical_devices('GPU')))\n"
115         )
116         try:
117             res = self._wsl_bash(cmd)
118             except FileNotFoundError:
119                 return False, "`wsl` not found ? is this running on Windows with WSL2
120 installed?"
121             except subprocess.TimeoutExpired:
122                 return False, "WSL healthcheck timed out."
123             out = (res.stdout or "").strip()
124             if res.returncode != 0:
125                 return False, f"WSL env check failed: {(res.stderr or out).strip()}"
126             return True, out
127
128     # ----- inference
129
130     def run_predict(self, video_win_path: str, output_win_dir: str) -> Optional[str]:
131         """Run predict.py on a video. Returns the Windows path to the produced CSV, or
132         None.
133
134         ``output_win_dir`` is a Windows directory (under the repo's output/) that
135         WSL writes into via its /mnt view, so the Windows process can read the CSV
136         back with no copy.
137         """
138         if not self.cfg.weights_path:
139             logger.error(
140                 "TrackNet weights_path is not set. Set indexing.tracknet.weights_path "
141                 "in config.json to the WSL path of your trained TrackNetV4 weights."
142             )
143             return None
144
145         os.makedirs(output_win_dir, exist_ok=True)
146         # Normalize separators so basename/splitext work when tests (or collector)
147         # pass Windows-style paths on a Linux host.
148         video_win_path = str(video_win_path).replace("\\", "/")
149         base = os.path.basename(video_win_path)
150         stem, ext = os.path.splitext(base)
151
152         # predict.py only accepts .avi/.mp4 and names the CSV "<stem>_predict.csv".
153         if ext.lower() not in (".mp4", ".avi"):
154             logger.error("TrackNet predict.py only supports .mp4/.avi (got %s)", ext)
155             return None
156
157         wsl_out = to_wsl_mnt_path(output_win_dir)
158         expected_csv_win = os.path.join(output_win_dir, f"{stem}_predict.csv")
159
160         # Resolve the video path WSL will read.
161         if self.cfg.stage_in_wsl:
162             staged = self._stage_video(video_win_path, base)
163             if staged is None:
164                 return None

```

```

162         wsl_video = staged
163     else:
164         wsl_video = to_wsl_mnt_path(video_win_path)
165
166     cmd = (
167         f"conda activate {self.cfg.conda_env} && "
168         f"cd {self.cfg.repo_dir}/src && "
169         f"python predict.py "
170         f"--video_path {shlex.quote(wsl_video)} "
171         f"--model_weights {shlex.quote(self.cfg.weights_path)} "
172         f"--output_dir {shlex.quote(wsl_out)} "
173         f"--queue_length {self.cfg.queue_length}"
174     )
175     logger.info("Running TrackNet inference on %s ...", base)
176     try:
177         res = self._wsl_bash(cmd)
178     except subprocess.TimeoutExpired:
179         logger.error("TrackNet inference timed out after %ss.",
180 self.cfg.timeout_sec)
181         return None
182
183     if res.returncode != 0:
184         logger.error("TrackNet predict.py failed:\n%s", (res.stderr or
185 res.stdout)[-2000:])
186         return None
187
188     if not os.path.exists(expected_csv_win):
189         logger.error("TrackNet finished but expected CSV not found at %s",
190 expected_csv_win)
191         return None
192
193     logger.info("TrackNet trajectory CSV: %s", expected_csv_win)
194     return expected_csv_win
195
196 def _stage_video(self, video_win_path: str, base: str) -> Optional[str]:
197     """Copy the input video into the WSL filesystem (avoids slow /mnt/c reads).
198
199     Returns the WSL path to the staged copy, or None on failure.
200     """
201     src_mnt = to_wsl_mnt_path(video_win_path)
202     dest = f"{self.cfg.wsl_stage_dir}/{base}"
203     cmd = f"mkdir -p {self.cfg.wsl_stage_dir} && cp {shlex.quote(src_mnt)}
204 {shlex.quote(dest)} && echo OK"
205     try:
206         res = self._wsl_bash(cmd)
207     except subprocess.TimeoutExpired:
208         logger.error("Staging video into WSL timed out.")
209         return None
210
211     if res.returncode != 0 or "OK" not in (res.stdout or ""):
212         logger.error("Failed to stage video into WSL: %s", (res.stderr or
213 res.stdout).strip())
214         return None
215     return dest
216
217 # ----- CSV -> windows
218
219 @staticmethod
220 def parse_trajectory_csv(csv_win_path: str) -> List[TrajectoryPoint]:
221     """Parse a TrackNet predict CSV (columns: Frame,Visibility,X,Y)."""
222     points: List[TrajectoryPoint] = []
223     with open(csv_win_path, newline="") as f:
224         reader = csv.DictReader(f)
225         for row in reader:
226             try:
227                 points.append(TrajectoryPoint(
228                     frame=int(row["Frame"]),

```

```

222             visible=int(row["Visibility"]) == 1,
223             x=float(row["X"]),
224             y=float(row["Y"]),
225         ))
226     except (KeyError, ValueError):
227         continue
228     points.sort(key=lambda p: p.frame)
229     return points
230
231 @staticmethod
232 def trajectory_to_action_windows(
233     points: List[TrajectoryPoint],
234     fps: float,
235     frame_width: float,
236     chunk_start: float = 0.0,
237     velocity_thresh: float = 0.02,
238     merge_gap: float = 2.0,
239     min_window_duration: float = 2.0,
240     chunk_index: Optional[int] = None,
241 ) -> List[Dict[str, Any]]:
242     """Convert a shuttle trajectory into candidate rally windows.
243
244     A frame is "active" when the shuttle is visible AND its inter-frame
245     displacement (normalised by frame width, per second) exceeds
246     ``velocity_thresh`` ? i.e. the shuttle is in flight. Active frames within
247     ``merge_gap`` seconds of each other are grouped into one window; short
248     windows are dropped. Mirrors HybridYoloSegmenter's windowing so the dict
249     shape feeding the AI-handoff phase is identical.
250     """
251     if fps <= 0:
252         fps = 30.0
253     if frame_width <= 0:
254         frame_width = 1280.0
255
256     # Compute per-frame normalised velocity from the last *visible* point.
257     active = [] # (frame_idx, velocity)
258     prev: Optional[TrajectoryPoint] = None
259     for p in points:
260         if not p.visible:
261             prev = None # break the trajectory; absent shuttle = not in flight
262             continue
263         if prev is None:
264             dframes = p.frame - prev.frame
265             if dframes > 0:
266                 dist = ((p.x - prev.x) ** 2 + (p.y - prev.y) ** 2) ** 0.5
267                 velocity = (dist / frame_width) * (fps / dframes)
268                 if velocity > velocity_thresh:
269                     active.append((p.frame, velocity))
270             prev = p
271
272     if not active:
273         return []
274
275     max_gap_frames = int(merge_gap * fps)
276
277     def _finalize(group: List[Tuple[int, float]]) -> Dict[str, Any]:
278         vels = [v for _, v in group]
279         return {
280             "start": group[0][0] / fps + chunk_start,
281             "end": group[-1][0] / fps + chunk_start,
282             "active_frame_count": len(group),
283             "peak_velocity": max(vels),
284             "mean_velocity": sum(vels) / len(vels),
285             "track_id_count": 1, # single shuttle; kept for window-shape parity
286             "chunk_index": chunk_index,

```

```

287         }
288
289         windows = []
290         group = [active[0]]
291         for af in active[1:]:
292             if af[0] - group[-1][0] <= max_gap_frames:
293                 group.append(af)
294             else:
295                 windows.append(_finalize(group))
296                 group = [af]
297         windows.append(_finalize(group))
298
299         return [w for w in windows if (w["end"] - w["start"]) >= min_window_duration]
300

```

22. user (2026-06-10T06:03:00.702Z)

```

1     """Abstract base for detector runners (TrackNet, WASB, future native, etc.).
2
3     This completes the Detector Runner Abstraction (low-regret item 6 continuation / item 1
4     of the post-observability re-prioritization) to de-risk the WSL/conda dependency for
5     the high-accuracy shuttle trajectory substrate used by tracknet_hybrid / wasb_hybrid.
6
7     The ABC lets us:
8     - Keep the current WSL adapters working unchanged for Windows users.
9     - Swap in (or A/B test) a Linux-native implementation later with zero changes to
10    the hybrid segmenters or the model-agnostic windowing brain.
11    - Provide a single place to evolve the healthcheck / error contract and thread
12    detector readiness into per-video reports.
13
14    Contract (deliberately matches what the concrete runners + hybrids + tests already did):
15    - run_predict(video_win_path: str, output_win_dir: str) -> Optional[str]
16      Run detector, return the Windows-side path to the produced trajectory CSV
17      (e.g. ".../clip_predict.csv" or ".../clip_wasb.csv"), or None on any failure.
18      The caller (hybrid) is responsible for output dir; impls may stage the video
19      into WSL fs for perf.
20
21    - healthcheck() -> tuple[bool, str]
22      Return (ok, message). Never raises. "ok" decides whether to proceed to run_predict.
23      Message is logged (and can be surfaced to reports or UX in future).
24      Intended to be cheap (env existence, import, GPU visibility, weights present).
25
26    Existing behavior is 100% preserved. The only additions are the inheritance and
27    the explicit common type for future implementers.
28
29    How to add a new runner (for docs / future contributors):
30    1. Subclass DetectorRunner in a new or existing module under pipeline/detectors/.
31    2. Implement run_predict and healthcheck (return (bool, str) tuple).
32    3. If it produces the same CSV schema (Frame,Visibility,X,Y), reuse or re-export
33       TrackNetRunner.parse_trajectory_csv and .trajectory_to_action_windows (they are
34       deliberately static and model-agnostic). If the new detector needs different
35       windowing, it can implement its own or we generalize later.
36    4. Lazy-import heavy native deps inside the methods (so importing the segmenter
37       registry never pulls torch/tf on a machine that doesn't have them).
38    5. Update the two hybrid segmenters (or introduce a detector= choice) to accept
39       DetectorRunner instances (they already do the right thing via duck-typing + the
40       type hints we add here).
41    6. Add unit tests that exercise the contract (isinstance, healthcheck shape, run
42       returns str|None) using mocks for the external bits ? see test_tracknet_runner.py.
43    7. No change to eval/calibration code paths (they consume the CSV + the static
44       windowing helpers directly).
45
46    Optional Linux-native path (the main de-risk goal):
47    A LinuxNativeRunner (or ONNXRunner) would:
48    - Skip all wsl / _stage / conda activate.

```

```

49 - Load weights locally (torch or onnx) or call a native binary.
50 - Write the exact same CSV format so trajectory_to_action_windows works verbatim.
51 - healthcheck can just check "weights file exists and torch importable + CUDA?".
52 This removes the WSL requirement for Linux users / future SaaS workers while
53 keeping the exact same candidate quality characteristics (or better, once native
54 weights are available).
55
56 Error surfacing into reports:
57 The healthcheck message is already logged by the hybrids on failure. The reporting
58 harvest sees the outcome as "0 candidates" + detector name. When we wire richer
59 segmenter Results (or pass detector_health into emit_indexer_report context), the
60 (ok, msg) can be recorded explicitly under segmentation.details or a new detector
61 stage. The ABC makes that future change local to one place.
62
63 See also:
64 - backend/pipeline/segmenters/{tracknet_hybrid,wasb_hybrid}.py (the two callers)
65 - tests/test_tracknet_runner.py (contract + pure logic tests; hybrids mock the runner)
66 - docs/TRACKNET_WSL_SETUP.md and CODE_MAP.md (2.2 Detectors section)
67 ""
68
69 from abc import ABC, abstractmethod
70 from typing import Optional, Tuple
71
72
73 class DetectorRunner(ABC):
74     """Common interface for shuttle/ball trajectory detectors (WSL or otherwise)."""
75
76     @abstractmethod
77     def run_predict(self, video_win_path: str, output_win_dir: str) -> Optional[str]:
78         """Run inference and return path to trajectory CSV (or None on failure)."""
79         pass
80
81     @abstractmethod
82     def healthcheck(self) -> Tuple[bool, str]:
83         """Quick check if the runner (env, weights, etc.) is usable.
84
85         Returns (ok, message). Never raises. Cheap to call before a long run.
86         """
87         pass
88

```

23. user (2026-06-10T06:03:10.869Z)

```

1     """WASB shuttle-trajectory inference on an arbitrary frame sequence (runs in WSL).
2
3     WASB-SBDT only ships *dataset* evaluation; this is a thin wrapper that reuses its
4     own building blocks (ImageDataset transform + detector + online tracker) to run on
5     any folder of extracted frames and emit a trajectory CSV in the exact shape the
6     indexer's `tracknet_runner.parse_trajectory_csv` already consumes:
7
8         Frame,Visibility,X,Y
9
10    It must run INSIDE the WSL `wasb` conda env, from the WASB-SBDT `src/` dir (it
11    imports WASB's `detectors`/`trackers`/`dataloaders`). The Windows-side adapter
12    (`wasb_runner.py`) stages this file + the frames into WSL and invokes it.
13
14    Resumable across reboots
15    -----
16    The GPU detector pass (one forward per sliding window ? ~21k for a 6-min clip) is
17    the slow, expensive part; the CPU tracker pass that turns detections into a
18    trajectory is cheap and *stateful* (it must see every frame in order, so it can't
19    resume mid-stream). We exploit that split:
20
21    * The detector pass writes its per-window raw outputs incrementally to
22    ``<cache>/det_raw.jsonl`` (flushed+fsync'd every ``--flush-every`` windows) and

```

```

23     records progress in ``<cache>/manifest.json`` (atomic write). A reboot loses at
24     most ``--flush-every`` windows of GPU work instead of the whole pass.
25     * On restart we replay the cached windows, skip the DataLoader ahead to the first
26     un-cached window, and continue. Once every window is cached, the detector pass
27     is skipped entirely and we just re-run the (cheap, deterministic) tracker over
28     the full ordered cache -> trajectory CSV. So a lost trajectory CSV costs seconds,
29     not the GPU hour.
30
31     The cache is keyed by the output path (``<out>_wasbcache/`` by default), lives next
32     to the output (NOT in %TEMP%), and is invalidated automatically if the frames dir,
33     weights, sport, window size, or frame count change.
34
35     Usage (in WSL, from ~/models/WASB-SBDT/src):
36     python wasb_infer.py --frames_dir /path/frames --weights
37     ../pretrained_weights/wasb_badminton_best.pth.tar \
38     --out /path/traj.csv [--sport badminton] [--limit N] [--cache-dir DIR] [--flush-
39     every 1000] [--fresh]
40     """
41     import os
42     import os.path as osp
43     import csv
44     import glob
45     import json
46     import argparse
47     import logging
48     from collections import defaultdict
49
50     logger = logging.getLogger(__name__)
51
52     CACHE_VERSION = 1
53     _DET_FILE = "det_raw.jsonl"
54     _MANIFEST_FILE = "manifest.json"
55
56     def _build_cfg(weights: str, sport: str):
57         """Compose the WASB eval config via Hydra, overriding model/weights/device."""
58         from hydra import compose, initialize
59         with initialize(config_path="configs", version_base=None):
60             cfg = compose(config_name="eval", overrides=[
61                 f"dataset={sport}",
62                 "model=wasb",
63                 f"detector.model_path={weights}",
64                 "runner.device=cuda",
65                 "runner.gpus=[0]",          # this box has a single GPU; default config assumes
66             ])
67         return cfg
68
69     def _frame_id(path: str) -> int:
70         """Best-effort integer frame id from a frame filename (e.g. frame_000180.png ->
71         180)."""
72         stem = osp.splitext(osp.basename(path))[0]
73         digits = "".join(ch for ch in stem if ch.isdigit())
74         return int(digits) if digits else -1
75
76     # ----- #
77     # Resumable detector cache (pure-Python, no torch/numpy ? unit-testable)
78     # ----- #
79     def _cache_paths(cache_dir: str):
80         return osp.join(cache_dir, _DET_FILE), osp.join(cache_dir, _MANIFEST_FILE)
81
82
83     def default_cache_dir(out_csv: str) -> str:

```

```

84         """Stable cache dir next to the trajectory output (survives reboots; not %TEMP%)."""
85         d = osp.dirname(osp.abspath(out_csv))
86         stem = osp.splitext(osp.basename(out_csv))[0]
87         return osp.join(d, stem + "_wasbcache")
88
89
90     def _atomic_write_text(path: str, text: str) -> None:
91         """Write then rename so a crash mid-write can't leave a half-written file."""
92         tmp = path + ".tmp"
93         with open(tmp, "w") as f:
94             f.write(text)
95             f.flush()
96             os.fsync(f.fileno())
97         os.replace(tmp, path)
98
99
100    def _det_plain(det) -> list:
101        """A detector candidate dict {'xy': np.array([x,y]), 'score', 'scale'} -> JSON-able
102        list."""
103        xy = det["xy"]
104        return [float(xy[0]), float(xy[1]), float(det["score"]), int(det["scale"])]
105
106    def _dets_to_tracker(plain_list):
107        """Inverse of _det_plain: rebuild the dicts the tracker expects (xy MUST be a numpy
108        array,
109        the tracker does vector math / np.linalg.norm on it)."""
110        import numpy as np
111        return [{"xy": np.array([p[0], p[1]], dtype=float), "score": p[2], "scale": p[3]}
112                for p in plain_list]
113
114    def write_manifest(cache_dir: str, manifest: dict) -> None:
115        _, mpath = _cache_paths(cache_dir)
116        _atomic_write_text(mpath, json.dumps(manifest, indent=2))
117
118
119    def read_manifest(cache_dir: str):
120        _, mpath = _cache_paths(cache_dir)
121        if not osp.exists(mpath):
122            return None
123        try:
124            with open(mpath) as f:
125                return json.load(f)
126        except (json.JSONDecodeError, OSError):
127            return None
128
129
130    def manifest_compatible(manifest: dict, *, frames_dir: str, weights: str, sport: str,
131                           frames_in: int, frames_total: int) -> bool:
132        """A cache is reusable only if the run that produced it matches this run's
133        inputs."""
134        if not manifest or manifest.get("version") != CACHE_VERSION:
135            return False
136        return (
137            manifest.get("frames_dir") == osp.abspath(frames_dir)
138            and manifest.get("weights") == weights
139            and manifest.get("sport") == sport
140            and manifest.get("frames_in") == frames_in
141            and manifest.get("frames_total") == frames_total
142        )
143
144    def reset_cache(cache_dir: str) -> None:
145        """Drop any existing cache so the next run starts clean."""

```



```

146     det_jsonl, mpath = _cache_paths(cache_dir)
147     for p in (det_jsonl, mpath):
148         if osp.exists(p):
149             os.remove(p)
150
151
152 def load_and_clean_cache(cache_dir: str):
153     """Read the longest *contiguous* (w == line index) valid prefix of det_raw.jsonl,
154     rebuild the per-frame detection accumulation, and rewrite the file to exactly that
155     prefix so a trailing half-written line from a crash can't corrupt future appends.
156
157     Returns (windows_done, det_by_fid) where det_by_fid maps frame_id -> list of
158     [x, y, score, scale] candidates (accumulated across every window the frame is in,
159     in window order ? identical to a non-resumed run).
160     """
161     det_jsonl, _ = _cache_paths(cache_dir)
162     det_by_fid = defaultdict(list)
163     valid = []
164     if osp.exists(det_jsonl):
165         with open(det_jsonl, "r") as f:
166             for line in f:
167                 s = line.strip()
168                 if not s:
169                     continue
170                 try:
171                     obj = json.loads(s)
172                 except json.JSONDecodeError:
173                     break # truncated trailing line (crash mid-
flush)
174                 if obj.get("w") != len(valid): # non-contiguous -> stop at the gap
175                     break
176                 valid.append(s)
177                 for fid, plain in obj["f"]:
178                     det_by_fid[int(fid)].extend([list(p) for p in plain])
179                 # Guarantee a clean append boundary by rewriting only the good prefix.
180                 _atomic_write_text(det_jsonl, ("\n".join(valid) + "\n") if valid else "")
181     return len(valid), det_by_fid
182
183
184 def _append_windows(fh, objs) -> None:
185     """Append window records as JSONL and durably flush (bounded loss on crash)."""
186     for o in objs:
187         fh.write(json.dumps(o, separators=(",", ":")) + "\n")
188     fh.flush()
189     os.fsync(fh.fileno())
190
191
192 def _atomic_write_trajectory(out_csv: str, rows) -> None:
193     os.makedirs(osp.dirname(osp.abspath(out_csv)), exist_ok=True)
194     tmp = out_csv + ".tmp"
195     with open(tmp, "w", newline="") as f:
196         w = csv.writer(f)
197         w.writerow(["Frame", "Visibility", "X", "Y"])
198         w.writerows(rows)
199         f.flush()
200         os.fsync(f.fileno())
201     os.replace(tmp, out_csv)
202
203
204 # ----- #
205 # Inference
206 # ----- #
207 def run(frames_dir: str, weights: str, out_csv: str, sport: str = "badminton",
208         limit: int = 0, batch_size: int = 8, num_workers: int = 4,
209         log_every_batches: int = 50, cache_dir: str = None,

```

```

210         flush_every: int = 1000, fresh: bool = False) -> str:
211     import time
212     import torch
213     from torch.utils.data import DataLoader
214     from detectors import build_detector
215     from trackers import build_tracker
216     from dataloaders import build_img_transforms, build_seq_transforms
217     from dataloaders.dataset_loader import ImageDataset
218     from utils import Center
219
220     cfg = _build_cfg(weights, sport)
221     frames_in = int(cfg["model"]["frames_in"])
222     input_wh = (int(cfg["model"]["inp_width"]), int(cfg["model"]["inp_height"]))
223     output_wh = (int(cfg["model"]["out_width"]), int(cfg["model"]["out_height"]))
224
225     frames = sorted(glob.glob(osp.join(frames_dir, "*.png"))) +
glob.glob(osp.join(frames_dir, "*.jpg")))
226     if limit:
227         frames = frames[:limit]
228     if len(frames) < frames_in:
229         raise SystemExit(f"need >= {frames_in} frames, found {len(frames)} in
{frames_dir}")
230
231     # Sliding windows of `frames_in` consecutive frames (the unit of GPU work +
checkpoint).
232     samples = []
233     for i in range(len(frames) - frames_in + 1):
234         window = frames[i:i + frames_in]
235         annos = [{"frame_path": p, "center": Center(is_visible=False, x=-1.0, y=-1.0)}
for p in window]
236         samples.append({"frames": window, "annos": annos})
237     windows_total = len(samples)
238
239     # ---- cache setup / resume decision ----- #
240     if cache_dir is None:
241         cache_dir = default_cache_dir(out_csv)
242         os.makedirs(cache_dir, exist_ok=True)
243         det_jsonl, _ = _cache_paths(cache_dir)
244         manifest = None if fresh else read_manifest(cache_dir)
245         resume = manifest_compatible(
246             manifest or {}, frames_dir=frames_dir, weights=weights, sport=sport,
247             frames_in=frames_in, frames_total=len(frames),
248         )
249     if resume:
250         windows_done, det_by_fid = load_and_clean_cache(cache_dir)
251         logger.info(f"resuming from cache: {windows_done}/{windows_total} windows
already detected")
252     else:
253         if manifest is not None:
254             logger.info("cache incompatible with this run -> starting fresh")
255             reset_cache(cache_dir)
256             windows_done, det_by_fid = 0, defaultdict(list)
257
258     base_manifest = {
259         "version": CACHE_VERSION,
260         "frames_dir": osp.abspath(frames_dir),
261         "weights": weights,
262         "sport": sport,
263         "frames_in": frames_in,
264         "frames_total": len(frames),
265         "windows_total": windows_total,
266         "windows_done": windows_done,
267     }
268     write_manifest(cache_dir, base_manifest)
269

```

```

270 # ---- detector pass over the un-cached windows ----- #
271 remaining = samples[windows_done:]
272 if remaining:
273     _, transform_test = build_img_transforms(cfg)
274     try:
275         _, seq_transform_test = build_seq_transforms(cfg)
276     except Exception:
277         seq_transform_test = None
278     ds = ImageDataset(cfg, remaining, input_wh, output_wh,
279                      transform=transform_test, seq_transform=seq_transform_test,
is_train=False)
280     loader = DataLoader(ds, batch_size=batch_size, shuffle=False,
num_workers=num_workers)
281     detector = build_detector(cfg)
282
283     t0 = time.time()
284     done = windows_done
285     win_idx = windows_done
286     log_every = max(1, log_every_batches)
287     flush_n = max(1, flush_every)
288     pending = []
289     logger.info(f"detecting on {len(remaining)} remaining windows "
290               f"(total {windows_total}, batch={batch_size},
workers={num_workers})...")
291     det_f = open(det_jsonl, "a")
292     try:
293         with torch.no_grad():
294             for bi, batch in enumerate(loader):
295                 imgs, _hms, trans = batch[0], batch[1], batch[2]
296                 img_paths = [list(t) for t in batch[-1]] # [frame_in][batch] ->
path
297
298                 batch_results, _ = detector.run_tensor(imgs, trans)
299                 nb = imgs.shape[0]
300                 for ib in range(nb):
301                     frames_payload = []
302                     for ie in sorted(batch_results[ib].keys()):
303                         p = img_paths[ie][ib]
304                         fid = _frame_id(p)
305                         plain = [_det_plain(d) for d in batch_results[ib][ie]]
306                         frames_payload.append([fid, plain])
307                         det_by_fid[fid].extend(plain)
308                     pending.append({"w": win_idx, "f": frames_payload})
309                     win_idx += 1
310                 done += nb
311                 if len(pending) >= flush_n or done >= windows_total:
312                     _append_windows(det_f, pending)
313                     pending = []
314                     base_manifest["windows_done"] = done
315                     write_manifest(cache_dir, base_manifest)
316                 if (bi + 1) % log_every == 0 or done >= windows_total:
317                     el = time.time() - t0
318                     new_done = done - windows_done
319                     rate = new_done / el if el > 0 else 0.0
320                     eta = (windows_total - done) / rate if rate > 0 else 0.0
321                     logger.info(f"{done}/{windows_total} "
322                               f"({100 * done // max(1, windows_total)}%) |
{rate:.1f} win/s | "
323                               f"elapsed {el:.0f}s | ETA {eta:.0f}s")
324                 if pending:
325                     _append_windows(det_f, pending)
326                     base_manifest["windows_done"] = done
327                     write_manifest(cache_dir, base_manifest)
328             finally:
329                 det_f.close()
330                 logger.info(f"detection done in {time.time() - t0:.0f}s; running tracker...")

```

```

330     else:
331         logger.info("detector cache already complete -> tracker-only re-run")
332
333     # ---- tracker re-run over the full ordered cache (cheap, deterministic) ---- #
334     tracker = build_tracker(cfg)
335     tracker.refresh()
336     rows = []
337     for fid in sorted(det_by_fid.keys()):
338         r = tracker.update(_dets_to_tracker(det_by_fid[fid]))
339         rows.append((fid, 1 if r.get("visi") else 0, r.get("x", -1), r.get("y", -1)))
340
341     _atomic_write_trajectory(out_csv, rows)
342     visible = sum(1 for _, v, *_ in rows if v)
343     logger.info(f"wrote {len(rows)} rows -> {out_csv}")
344     logger.info(f"visible-shuttle frames: {visible}/{len(rows)}")
345     return out_csv
346
347
348 def extract_frames(video_path: str, frames_dir: str) -> str:
349     """Decode a video to PNG frames named by source frame index (000000000.png ...).
350
351     Idempotent: if the dir already holds the expected number of frames (matching the
352     container's reported frame count) we skip re-decoding ? re-extraction after a
353     reboot is wasteful. NOTE: cv2 PNG extraction is slow; for big clips prefer
354     extracting with ffmpeg on the host and passing --frames_dir.
355     """
356     import cv2
357     os.makedirs(frames_dir, exist_ok=True)
358     cap = cv2.VideoCapture(video_path)
359     if not cap.isOpened():
360         raise SystemExit(f"cannot open video: {video_path}")
361     expected = int(cap.get(cv2.CAP_PROP_FRAME_COUNT) or 0)
362     existing = len(glob.glob(osp.join(frames_dir, "*.png")))
363     if expected > 0 and existing >= expected:
364         cap.release()
365         logger.info(f"reusing {existing} already-extracted frames -> {frames_dir}")
366         return frames_dir
367     i = 0
368     while True:
369         ok, frame = cap.read()
370         if not ok:
371             break
372         cv2.imwrite(osp.join(frames_dir, f"{i:08d}.png"), frame)
373         i += 1
374     cap.release()
375     logger.info(f"extracted {i} frames -> {frames_dir}")
376     return frames_dir
377
378
379 def main():
380     ap = argparse.ArgumentParser()
381     ap.add_argument("--frames_dir", help="folder of pre-extracted frames")
382     ap.add_argument("--video", help="video file; frames are extracted internally (in
WSL)")
383     ap.add_argument("--frames_out_dir", help="where to extract frames when --video is
used")
384     ap.add_argument("--weights", required=True)
385     ap.add_argument("--out", required=True)
386     ap.add_argument("--sport", default="badminton")
387     ap.add_argument("--limit", type=int, default=0, help="cap #frames (for quick
tests)")
388     ap.add_argument("--batch-size", type=int, default=8)
389     ap.add_argument("--num-workers", type=int, default=4)
390     ap.add_argument("--log-every-batches", type=int, default=50, help="progress log
cadence")

```

```

391         ap.add_argument("--cache-dir", default=None,
392                         help="resumable detector cache dir (default: <out>_wasbcache next to
--out)")
393         ap.add_argument("--flush-every", type=int, default=1000,
394                         help="fsync the detector cache every N windows (bounds reboot
loss)")
395         ap.add_argument("--fresh", action="store_true", help="ignore any existing cache and
start over")
396         args = ap.parse_args()
397
398         frames_dir = args.frames_dir
399         if args.video:
400             frames_dir = extract_frames(args.video, args.frames_out_dir or (args.video +
"_frames"))
401         if not frames_dir:
402             raise SystemExit("provide --frames_dir or --video")
403         run(frames_dir, args.weights, args.out, args.sport, args.limit,
404             batch_size=args.batch_size, num_workers=args.num_workers,
405             log_every_batches=args.log_every_batches, cache_dir=args.cache_dir,
406             flush_every=args.flush_every, fresh=args.fresh)
407
408
409     if __name__ == "__main__":
410         main()
411

```

24. user (2026-06-10T06:03:33.625Z)

```

1     """Rally-start gate: suppress between-point false fires using shuttle net-crossings.
2
3     WASB (shuttle-only) over-detects: when a player flicks/tosses the shuttle back to
4     the opponent for service after losing a point, that single trajectory looks like a
5     rally. The discriminator is structural, not duration-based: a real rally has the
6     shuttle crossing the net MULTIPLE times (each shot lands on the other side); a
7     single service feed crosses ~once.
8
9     This module is camera-agnostic and needs NO new model ? it works on the trajectory
10    we already produce (Frame,Visibility,X,Y). It estimates the play axis (the image
11    axis with the larger shuttle spread) and the net position (median shuttle coord on
12    that axis ? the shuttle clusters at the two ends and crosses the net region often,
13    so the median sits near the net), then counts sign changes of (coord - net) across
14    the visible points inside a candidate window, with a deadband to ignore jitter.
15
16    Gate A in the over-fire fix spectrum (B = player-stance state machine, future).
17    Pure functions only ? no I/O, unit-tested; intended as a post-filter on the
18    windows from trajectory_to_action_windows, or to fold into it later.
19    """
20    from __future__ import annotations
21
22    from dataclasses import dataclass
23    from typing import List, Sequence, Tuple
24
25    Interval = Tuple[float, float]
26
27
28    @dataclass
29    class GatedWindow:
30        start: float
31        end: float
32        crossings: int
33        visible_points: int
34        kept: bool
35
36
37    def _spread(vals: Sequence[float]) -> float:

```

```

38     """Robust spread = p95 - p5 (ignores a few outlier detections)."""
39     if len(vals) < 2:
40         return 0.0
41     s = sorted(vals)
42     lo = s[max(0, int(0.05 * (len(s) - 1)))]
43     hi = s[min(len(s) - 1, int(0.95 * (len(s) - 1)))]
44     return hi - lo
45
46
47 def _median(vals: Sequence[float]) -> float:
48     if not vals:
49         return 0.0
50     s = sorted(vals)
51     n = len(s)
52     return s[n // 2] if n % 2 else 0.5 * (s[n // 2 - 1] + s[n // 2])
53
54
55 def estimate_play_axis(points) -> Tuple[str, float, float]:
56     """Return (axis 'x'/'y', net_value, span) from visible points.
57
58     Play axis = the image axis along which the shuttle travels most (larger spread).
59     For a baseline camera that's Y (players top/bottom); for a side camera, X.
60     """
61     xs = [p.x for p in points if p.visible]
62     ys = [p.y for p in points if p.visible]
63     sx, sy = _spread(xs), _spread(ys)
64     if sx >= sy:
65         return "x", _median(xs), sx
66     return "y", _median(ys), sy
67
68
69 def count_net_crossings(window_points, axis: str, net: float, deadband: float) -> int:
70     """Count sign changes of (coord - net) across visible points, ignoring the
71     deadband zone near the net (so jitter at the net doesn't inflate the count)."""
72     crossings = 0
73     last_side = None
74     for p in window_points:
75         if not p.visible:
76             continue
77         v = (p.x if axis == "x" else p.y) - net
78         if abs(v) < deadband:
79             continue
80         side = v > 0
81         if last_side is not None and side != last_side:
82             crossings += 1
83         last_side = side
84     return crossings
85
86
87 def _points_in(points, fps: float, start: float, end: float):
88     return [p for p in points if start * fps <= p.frame <= end * fps]
89
90
91 def gate_windows(points, windows: List[Interval], fps: float,
92                 min_crossings: int = 2, deadband_frac: float = 0.06) ->
List[GatedWindow]:
93     """Annotate each (start,end) window with its shuttle net-crossing count and a
94     keep decision (crossings >= min_crossings). Deadband is a fraction of the play
95     axis span. Net/axis are estimated once over the whole trajectory (stable)."""
96     axis, net, span = estimate_play_axis(points)
97     deadband = deadband_frac * span
98     out: List[GatedWindow] = []
99     for (s, e) in windows:
100         wp = _points_in(points, fps, s, e)
101         c = count_net_crossings(wp, axis, net, deadband)

```

```

102         vis = sum(1 for p in wp if p.visible)
103         out.append(GatedWindow(start=s, end=e, crossings=c, visible_points=vis,
104                               kept=c >= min_crossings))
105     return out
106
107
108     def apply_gate(points, windows: List[Interval], fps: float,
109                   min_crossings: int = 2, deadband_frac: float = 0.06) -> List[Interval]:
110         """Convenience: return only the windows that pass the crossing gate."""
111         return [(g.start, g.end) for g in
112                 gate_windows(points, windows, fps, min_crossings, deadband_frac) if g.kept]
113

```

25. user (2026-06-10T06:03:34.500Z)

```

1     from abc import ABC, abstractmethod
2     from pydantic import BaseModel
3     from typing import List, Dict, Any, Optional
4
5     class ValidationResult(BaseModel):
6         validator_name: str
7         passed: bool
8         message: str
9         details: Dict[str, Any] = {}
10
11     class VideoValidator(ABC):
12         @property
13         @abstractmethod
14         def name(self) -> str:
15             """Unique identifier for this validator."""
16             pass
17
18         @property
19         @abstractmethod
20         def description(self) -> str:
21             """Human-readable explanation of what this validator checks."""
22             pass
23
24         @abstractmethod
25         def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
26             """
27             Executes validation.
28             - video_path: Local path to the MP4 file.
29             - config: Dict parsed from config.json.
30             - context: Dynamic context dict shared across validators (stores shared values
like ffprobe metadata).
31             """
32             pass
33
34     class ValidationRegistry:
35         def __init__(self):
36             self._validators: List[VideoValidator] = []
37
38         def register(self, validator: VideoValidator):
39             """Registers a new validator."""
40             self._validators.append(validator)
41
42         def run_all_with_context(self, video_path: str, config: Dict[str, Any]):
43             """Runs all validators and returns ``(results, context)``.
44
45             The per-run ``context`` accumulates shared values (notably
46             ``context['ffprobe_meta']`` produced by FormatValidator). ``run_all``
47             discards it; callers that want media metadata for an observability report
48             use this variant instead of re-probing.

```

```

49         """
50         results = []
51         context = {} # Shared context dictionary (e.g. for caching ffprobe outputs)
52         for val in self._validators:
53             try:
54                 res = val.validate(video_path, config, context)
55                 results.append(res)
56             except Exception as e:
57                 results.append(ValidationResult(
58                     validator_name=val.name,
59                     passed=False,
60                     message=f"Validator threw an unexpected error: {str(e)}"
61                 ))
62         return results, context
63
64     def run_all(self, video_path: str, config: Dict[str, Any]) ->
List[ValidationResult]:
65         """Runs all registered validators in sequence, accumulating results."""
66         results, _ = self.run_all_with_context(video_path, config)
67         return results
68
69     # Shared global registry
70     registry = ValidationRegistry()
71

```

26. user (2026-06-10T06:03:34.721Z)

```

1     from backend.pipeline.validators.base import registry, VideoValidator, ValidationResult
2     from backend.pipeline.validators.format import FormatValidator
3     from backend.pipeline.validators.resolution_fps import ResolutionFPSValidator
4     from backend.pipeline.validators.pts_continuity import PTSContinuityValidator
5     from backend.pipeline.validators.scene_cut import SceneCutDensityValidator
6     from backend.pipeline.validators.blur import BlurValidator
7     from backend.pipeline.validators.relevance import RelevanceValidator
8
9     # Register all default validators in sequence (dependencies first)
10    registry.register(FormatValidator())
11    registry.register(ResolutionFPSValidator())
12    registry.register(PTSContinuityValidator())
13    registry.register(SceneCutDensityValidator())
14    registry.register(BlurValidator())
15    registry.register(RelevanceValidator())
16
17    __all__ = ["registry", "VideoValidator", "ValidationResult"]
18

```

27. user (2026-06-10T06:03:44.023Z)

```

1     import subprocess
2     import json
3     import os
4     from backend.pipeline.validators.base import VideoValidator, ValidationResult
5     from typing import Dict, Any
6
7     class FormatValidator(VideoValidator):
8         @property
9         def name(self) -> str:
10             return "format_check"
11
12         @property
13         def description(self) -> str:
14             return "Verifies the video is a valid, uncorrupted MP4 container and extracts
metadata."
15
16         def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])

```



```

-> ValidationResult:
17         if not os.path.exists(video_path):
18             return ValidationResult(
19                 validator_name=self.name,
20                 passed=False,
21                 message=f"File not found: {video_path}"
22             )
23
24         cmd = [
25             "ffprobe",
26             "-v", "error",
27             "-select_streams", "v:0",
28             "-show_format",
29             "-show_streams",
30             "-of", "json",
31             video_path
32         ]
33
34         try:
35             result = subprocess.run(cmd, capture_output=True, text=True, check=True)
36             info = json.loads(result.stdout)
37
38             if not info.get("streams") or not info.get("format"):
39                 return ValidationResult(
40                     validator_name=self.name,
41                     passed=False,
42                     message="Failed to retrieve valid video streams from the file."
43                 )
44
45             # Store the metadata in context for subsequent validators to utilize
46             context["ffprobe_meta"] = info
47
48             stream = info["streams"][0]
49             fmt = info["format"]
50
51             details = {
52                 "format_name": fmt.get("format_name"),
53                 "duration": float(fmt.get("duration", 0)),
54                 "size_bytes": int(fmt.get("size", 0)),
55                 "codec_name": stream.get("codec_name"),
56             }
57
58             # Check if container name includes mp4
59             if "mp4" not in fmt.get("format_name", "").lower():
60                 return ValidationResult(
61                     validator_name=self.name,
62                     passed=False,
63                     message=f"Invalid format: Found '{fmt.get('format_name')}', but MP4
is required.",
64                     details=details
65                 )
66
67             return ValidationResult(
68                 validator_name=self.name,
69                 passed=True,
70                 message="Video format is valid MP4 container.",
71                 details=details
72             )
73
74         except subprocess.CalledProcessError as e:
75             return ValidationResult(
76                 validator_name=self.name,
77                 passed=False,
78                 message="Video stream is corrupted or unreadable by ffprobe.",
79                 details={"error": e.stderr}

```

```

80         )
81     except Exception as e:
82         return ValidationResult(
83             validator_name=self.name,
84             passed=False,
85             message=f"Unexpected error running ffprobe: {str(e)}"
86         )
87

```

28. user (2026-06-10T06:03:44.840Z)

```

1     from backend.pipeline.validators.base import VideoValidator, ValidationResult
2     from typing import Dict, Any
3
4     class ResolutionFPSValidator(VideoValidator):
5         @property
6         def name(self) -> str:
7             return "resolution_fps_check"
8
9         @property
10        def description(self) -> str:
11            return "Checks if the video meets minimum resolution (e.g. 720p) and frame rate (e.g. 30fps) requirements."
12
13        def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
14            meta = context.get("ffprobe_meta")
15            if not meta or not meta.get("streams"):
16                return ValidationResult(
17                    validator_name=self.name,
18                    passed=False,
19                    message="Video metadata is missing. Run format_check validator first."
20                )
21
22            stream = meta["streams"][0]
23            width = int(stream.get("width", 0))
24            height = int(stream.get("height", 0))
25
26            # Parse frame rate (e.g. "30/1" or "30000/1001")
27            fps_str = stream.get("r_frame_rate", "0/1")
28            try:
29                num, den = map(int, fps_str.split("/"))
30                fps = num / den if den != 0 else 0
31            except ValueError:
32                fps = 0.0
33
34            details = {
35                "width": width,
36                "height": height,
37                "fps": round(fps, 2)
38            }
39
40            # Check bounds from config
41            val_cfg = config.get("validation", {})
42            min_w = val_cfg.get("min_resolution_width", 1280)
43            min_h = val_cfg.get("min_resolution_height", 720)
44            min_fps = val_cfg.get("min_fps", 29.9)
45
46            if width < min_w or height < min_h:
47                return ValidationResult(
48                    validator_name=self.name,
49                    passed=False,
50                    message=f"Low resolution: Found {width}x{height}, but minimum is
{min_w}x{min_h}.",
51                    details=details

```

```

52         )
53
54     if fps < min_fps:
55         return ValidationResult(
56             validator_name=self.name,
57             passed=False,
58             message=f"Low frame rate: Found {round(fps, 2)} FPS, but minimum is
{min_fps} FPS.",
59             details=details
60         )
61
62     return ValidationResult(
63         validator_name=self.name,
64         passed=True,
65         message="Video resolution and frame rate meet minimum standards.",
66         details=details
67     )
68

```

29. user (2026-06-10T06:03:45.081Z)

```

1     import subprocess
2     from backend.pipeline.validators.base import VideoValidator, ValidationResult
3     from typing import Dict, Any
4
5     class PTSContinuityValidator(VideoValidator):
6         @property
7         def name(self) -> str:
8             return "pts_continuity_check"
9
10        @property
11        def description(self) -> str:
12            return "Analyzes Presentation Time Stamps (PTS) of keyframes to check for
timeline gaps, jumps, or editing splices."
13
14        def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
15            # Command to extract packet PTS and flags (extremely fast, works on H.264 &
H.265)
16            cmd = [
17                "ffprobe",
18                "-v", "error",
19                "-select_streams", "v:0",
20                "-show_entries", "packet=pts_time,flags",
21                "-of", "csv=p=0",
22                video_path
23            ]
24
25            try:
26                result = subprocess.run(cmd, capture_output=True, text=True, check=True)
27                output = result.stdout.strip()
28
29                if not output:
30                    return ValidationResult(
31                        validator_name=self.name,
32                        passed=False,
33                        message="No packet timestamps could be extracted from the video
stream."
34                    )
35
36                timestamps = []
37                for line in output.split("\n"):
38                    line = line.strip()
39                    if line:
40                        parts = line.split(",")

```

```

41         if len(parts) == 2:
42             pts_str, flags = parts
43             if "K" in flags: # Check if keyframe
44                 try:
45                     timestamps.append(float(pts_str))
46                 except ValueError:
47                     continue
48
49     if len(timestamps) < 2:
50         return ValidationResult(
51             validator_name=self.name,
52             passed=True,
53             message="Video is too short to perform a PTS continuity check.",
54             details={"keyframes_count": len(timestamps)}
55         )
56
57     # Analyze gaps
58     max_gap = 0.0
59     backward_jumps = 0
60     gaps = []
61
62     for i in range(1, len(timestamps)):
63         gap = timestamps[i] - timestamps[i-1]
64         gaps.append(gap)
65         if gap < 0:
66             backward_jumps += 1
67         if gap > max_gap:
68             max_gap = gap
69
70     # An continuous GoPro recording usually has keyframe intervals (GOP size) of
71     1.0 to 4.0 seconds.
72     # If there's a keyframe gap larger than 8.0 seconds or a negative gap
73     (backward jump), it's highly likely spliced.
74     details = {
75         "keyframes_count": len(timestamps),
76         "max_keyframe_gap_seconds": round(max_gap, 2),
77         "backward_jumps_detected": backward_jumps,
78         "first_keyframe": timestamps[0],
79         "last_keyframe": timestamps[-1]
80     }
81
82     if backward_jumps > 0:
83         return ValidationResult(
84             validator_name=self.name,
85             passed=False,
86             message="Video integrity compromise: Backward timeline jumps
87             (splicing) detected.",
88             details=details
89         )
90
91     # Threshold for sudden gap (e.g. 10.0 seconds)
92     if max_gap > 10.0:
93         return ValidationResult(
94             validator_name=self.name,
95             passed=False,
96             message=f"Video integrity compromise: Detected a timestamp gap of
97             {round(max_gap, 2)}s, indicating edited or missing content.",
98             details=details
99         )
100
101     return ValidationResult(
102         validator_name=self.name,
103         passed=True,
104         message="PTS timeline is continuous with no abnormal gaps or backward
105         jumps.",

```

```

101         details=details
102     )
103
104     except subprocess.CalledProcessError as e:
105         return ValidationResult(
106             validator_name=self.name,
107             passed=False,
108             message="Failed to execute ffprobe PTS analysis.",
109             details={"error": e.stderr}
110         )
111     except Exception as e:
112         return ValidationResult(
113             validator_name=self.name,
114             passed=False,
115             message=f"Unexpected error in PTS validation: {str(e)}"
116         )
117

```

30. user (2026-06-10T06:03:52.533Z)

```

1     import cv2
2     import numpy as np
3     from backend.pipeline.validators.base import VideoValidator, ValidationResult
4     from typing import Dict, Any
5
6     class SceneCutDensityValidator(VideoValidator):
7         @property
8         def name(self) -> str:
9             return "scene_cut_check"
10
11        @property
12        def description(self) -> str:
13            return "Analyzes visual frames to detect abrupt scene cuts, which indicate an
edited or compiled video."
14
15        def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
16            cap = cv2.VideoCapture(video_path)
17            if not cap.isOpened():
18                return ValidationResult(
19                    validator_name=self.name,
20                    passed=False,
21                    message="Failed to open video file using OpenCV."
22                )
23
24            fps = cap.get(cv2.CAP_PROP_FPS)
25            total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
26            duration = total_frames / fps if fps > 0 else 0
27
28            if duration <= 0:
29                cap.release()
30                return ValidationResult(
31                    validator_name=self.name,
32                    passed=False,
33                    message="Invalid video duration."
34                )
35
36            # Sample at 2 frames per second to keep it fast
37            sample_interval = max(1, int(fps / 2))
38
39            cuts_detected = 0
40            prev_hist = None
41
42            frame_idx = 0
43            sampled_count = 0

```

```

44
45     # Adjust histogram comparison thresholds
46     # Using correlation: 1.0 is perfect match, < 0.6 indicates a sudden massive
scene change
47     cut_threshold = 0.6
48
49     while True:
50         cap.set(cv2.CAP_PROP_POS_FRAMES, frame_idx)
51         ret, frame = cap.read()
52         if not ret:
53             break
54
55         # Calculate color histogram
56         hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
57         hist = cv2.calcHist([hsv], [0, 1], None, [50, 60], [0, 180, 0, 256])
58         cv2.normalize(hist, hist, alpha=0, beta=1, norm_type=cv2.NORM_MINMAX)
59
60         if prev_hist is not None:
61             # Compare correlation
62             score = cv2.compareHist(prev_hist, hist, cv2.HISTCMP_CORREL)
63             if score < cut_threshold:
64                 cuts_detected += 1
65
66         prev_hist = hist
67         sampled_count += 1
68         frame_idx += sample_interval
69
70         # Cap check at 100 frames max (e.g. first 50 seconds or evenly spread) to
keep validation instant
71         if frame_idx >= total_frames or sampled_count > 100:
72             break
73
74         cap.release()
75
76         # Compute cuts per minute
77         sampled_duration = (sampled_count * sample_interval) / fps if fps > 0 else 1
78         cuts_per_minute = (cuts_detected / sampled_duration) * 60 if sampled_duration >
0 else 0
79
80         details = {
81             "cuts_detected": cuts_detected,
82             "sampled_frames": sampled_count,
83             "cuts_per_minute": round(cuts_per_minute, 3)
84         }
85
86         val_cfg = config.get("validation", {})
87         max_cuts_rate = val_cfg.get("max_scene_cuts_per_minute", 0.5)
88
89         if cuts_per_minute > max_cuts_rate:
90             return ValidationResult(
91                 validator_name=self.name,
92                 passed=False,
93                 message=f"Video splicing compromise: Detected {cuts_detected} cuts
({round(cuts_per_minute, 2)} cuts/min), which exceeds the continuous raw file limit of
{max_cuts_rate} cuts/min.",
94                 details=details
95             )
96
97         return ValidationResult(
98             validator_name=self.name,
99             passed=True,
100             message="Scene continuity check passed. No abnormal visual cuts detected.",
101             details=details
102         )
103

```

31. user (2026-06-10T06:03:53.006Z)

```
1      import cv2
2      from backend.pipeline.validators.base import VideoValidator, ValidationResult
3      from typing import Dict, Any
4
5      class BlurValidator(VideoValidator):
6          @property
7          def name(self) -> str:
8              return "blur_check"
9
10         @property
11         def description(self) -> str:
12             return "Samples frames across the video and checks if the footage is excessively
blurry or out of focus."
13
14         def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
15             cap = cv2.VideoCapture(video_path)
16             if not cap.isOpened():
17                 return ValidationResult(
18                     validator_name=self.name,
19                     passed=False,
20                     message="Failed to open video file for blur analysis."
21                 )
22
23             total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
24             if total_frames <= 0:
25                 cap.release()
26                 return ValidationResult(
27                     validator_name=self.name,
28                     passed=False,
29                     message="Invalid frame count."
30                 )
31
32             # Sample 15 frames evenly distributed across the video
33             num_samples = 15
34             step = max(1, total_frames // num_samples)
35
36             variances = []
37
38             for i in range(num_samples):
39                 frame_idx = i * step
40                 cap.set(cv2.CAP_PROP_POS_FRAMES, frame_idx)
41                 ret, frame = cap.read()
42                 if not ret:
43                     continue
44
45                 gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
46                 # Calculate the Laplacian variance (standard focus metric)
47                 laplacian_var = cv2.Laplacian(gray, cv2.CV_64F).var()
48                 variances.append(laplacian_var)
49
50             cap.release()
51
52             if not variances:
53                 return ValidationResult(
54                     validator_name=self.name,
55                     passed=False,
56                     message="Could not read frames to perform blur verification."
57                 )
58
59             avg_variance = sum(variances) / len(variances)
60
61             details = {
```

```

62         "avg_laplacian_variance": round(avg_variance, 2),
63         "min_laplacian_variance": round(min(variances), 2),
64         "max_laplacian_variance": round(max(variances), 2),
65         "sampled_frames_count": len(variances)
66     }
67
68     val_cfg = config.get("validation", {})
69     min_blur_threshold = val_cfg.get("min_blur_threshold", 100.0)
70
71     if avg_variance < min_blur_threshold:
72         return ValidationResult(
73             validator_name=self.name,
74             passed=False,
75             message=f"Video quality rejected: Focus check failed. Average sharpness
score of {round(avg_variance, 1)} is below the minimum required threshold of
{min_blur_threshold}.",
76             details=details
77         )
78
79     return ValidationResult(
80         validator_name=self.name,
81         passed=True,
82         message="Video sharpness meets focus requirements.",
83         details=details
84     )
85

```

32. user (2026-06-10T06:03:53.396Z)

```

1     import cv2
2     import numpy as np
3     from backend.pipeline.validators.base import VideoValidator, ValidationResult
4     from typing import Dict, Any
5
6     class RelevanceValidator(VideoValidator):
7         @property
8         def name(self) -> str:
9             return "relevance_check"
10
11         @property
12         def description(self) -> str:
13             return "Checks if the video content is relevant to badminton by scanning for
court color palettes and geometric features."
14
15         def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
16             cap = cv2.VideoCapture(video_path)
17             if not cap.isOpened():
18                 return ValidationResult(
19                     validator_name=self.name,
20                     passed=False,
21                     message="Failed to open video file for relevance validation."
22                 )
23
24             total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
25             if total_frames <= 0:
26                 cap.release()
27                 return ValidationResult(
28                     validator_name=self.name,
29                     passed=False,
30                     message="Invalid frame count."
31                 )
32
33             # Sample 10 frames across the video
34             num_samples = 10

```



```

35         step = max(1, total_frames // num_samples)
36
37         green_ratios = []
38
39         # Get sport profile config
40         sport_name = config.get("sport", "badminton")
41         from backend.sports import sport_registry
42         try:
43             sport_profile = sport_registry.get(sport_name)
44             sport_rel_cfg = sport_profile.get_relevance_config()
45         except KeyError:
46             sport_rel_cfg = {}
47
48         val_cfg = config.get("validation", {})
49         rel_cfg = val_cfg.get("relevance", {})
50
51         # Read bounds from sport profile with fallback to user config/defaults
52         hsv_lower = np.array(rel_cfg.get("court_green_lower",
sport_rel_cfg.get("court_green_lower", [35, 40, 40])))
53         hsv_upper = np.array(rel_cfg.get("court_green_upper",
sport_rel_cfg.get("court_green_upper", [85, 255, 255])))
54         min_court_ratio = rel_cfg.get("court_pixels_min_ratio",
sport_rel_cfg.get("court_pixels_min_ratio", 0.05))
55
56         for i in range(num_samples):
57             frame_idx = i * step
58             cap.set(cv2.CAP_PROP_POS_FRAMES, frame_idx)
59             ret, frame = cap.read()
60             if not ret:
61                 continue
62
63             # Downscale frame for speed
64             h, w = frame.shape[:2]
65             small_frame = cv2.resize(frame, (320, 240))
66
67             # Convert to HSV color space
68             hsv = cv2.cvtColor(small_frame, cv2.COLOR_BGR2HSV)
69
70             # Mask court green pixels
71             mask = cv2.inRange(hsv, hsv_lower, hsv_upper)
72             green_pixels = cv2.countNonZero(mask)
73             total_pixels = mask.shape[0] * mask.shape[1]
74
75             ratio = green_pixels / total_pixels
76             green_ratios.append(ratio)
77
78         cap.release()
79
80         if not green_ratios:
81             return ValidationResult(
82                 validator_name=self.name,
83                 passed=False,
84                 message="Failed to extract frames for relevance checking."
85             )
86
87         avg_green_ratio = sum(green_ratios) / len(green_ratios)
88
89         details = {
90             "avg_green_court_ratio": round(avg_green_ratio, 4),
91             "max_green_court_ratio": round(max(green_ratios), 4),
92             "min_green_court_ratio": round(min(green_ratios), 4),
93             "green_threshold": min_court_ratio
94         }
95
96         # If a certain amount of court color is visible, consider it relevant

```

```

97         if avg_green_ratio < min_court_ratio:
98             return ValidationResult(
99                 validator_name=self.name,
100                 passed=False,
101                 message=f"Video rejected: Relevance check failed. The scene does not
contain a recognizable {sport_name} court (court area found: {round(avg_green_ratio * 100, 2)}%
vs min required {round(min_court_ratio * 100, 2)}%).",
102                 details=details
103             )
104
105         return ValidationResult(
106             validator_name=self.name,
107             passed=True,
108             message=f"Video identified as relevant to {sport_name} ({sport_name} court
structures found).",
109             details=details
110         )
111

```

33. user (2026-06-10T06:04:10.577Z)

```

1     import os
2     import subprocess
3     import tempfile
4     import logging
5     from typing import List, Tuple, Union
6
7     logger = logging.getLogger(__name__)
8
9     class VideoCompiler:
10         def __init__(self, output_dir: str, end_padding: float = 1.0, begin_padding: float =
0.5):
11             self.output_dir = output_dir
12             self.end_padding = end_padding
13             self.begin_padding = begin_padding
14             os.makedirs(output_dir, exist_ok=True)
15
16         @staticmethod
17         def _parse_time(val: Union[int, float, str]) -> float:
18             """Normalizes a timestamp value to float seconds.
19             Handles: float/int (pass-through), plain numeric strings ("12.4"),
20             and colon-separated strings ("MM:SS", "HH:MM:SS", "H:MM:SS.ms").
21             This mirrors the parse_time logic used in gemini.py to ensure
22             the stitcher can always consume whatever format the indexer produced."""
23             if isinstance(val, (int, float)):
24                 return float(val)
25             if isinstance(val, str):
26                 val = val.strip()
27                 if ":" in val:
28                     # Handle MM:SS, HH:MM:SS, etc.
29                     parts = val.split(":")
30                     parts.reverse()
31                     total = 0.0
32                     for i, p in enumerate(parts):
33                         try:
34                             total += float(p) * (60 ** i)
35                         except ValueError:
36                             pass
37                     return total
38                 try:
39                     return float(val)
40                 except ValueError:
41                     logger.warning(f"Could not parse timestamp value '{val}' as a number.
Defaulting to 0.0.")
42             return 0.0

```

```

43         logger.warning(f"Unexpected timestamp type {type(val).__name__} for value
'{val}'. Defaulting to 0.0.")
44         return 0.0
45
46     def _get_video_duration(self, video_path: str) -> float:
47         """Probes the video file to get its total duration in seconds.
48         Returns 0.0 if it cannot be determined."""
49         try:
50             result = subprocess.run(
51                 [
52                     "ffprobe", "-v", "error",
53                     "-show_entries", "format=duration",
54                     "-of", "default=noprint_wrappers=1:nokey=1",
55                     video_path
56                 ],
57                 capture_output=True, text=True, check=True
58             )
59             return float(result.stdout.strip())
60         except Exception as e:
61             logger.warning(f"Could not determine video duration via ffprobe: {e}")
62             return 0.0
63
64     def compile_highlights(self, raw_video_path: str, segments: List[Tuple[float,
float]], output_filename: str) -> str:
65         """
66         Extracts selected segments from a raw video and stitches them together
67         into a seamless highlights file, applying end_padding to prevent abrupt endings.
68         """
69         if not segments:
70             raise ValueError("No highlight segments provided to compile.")
71
72         if not os.path.exists(raw_video_path):
73             raise FileNotFoundError(f"[COMPILER ERROR] Source video file not found:
{raw_video_path}")
74
75         output_path = os.path.join(self.output_dir, output_filename)
76
77         # Probe video duration so we can detect out-of-range segments
78         video_duration = self._get_video_duration(raw_video_path)
79         if video_duration > 0:
80             logger.info(f"Source video duration: {video_duration:.2f}s")
81         else:
82             logger.warning("Video duration unknown. Cannot validate segment boundaries
against video length.")
83
84         # We will create temporary clips and stitch them using FFmpeg concat demuxer
85         temp_clips = []
86         skipped_segments = []
87
88         # Create a temp directory within output_dir
89         with tempfile.TemporaryDirectory(dir=self.output_dir) as tmp_dir:
90             logger.info(f"Trimming {len(segments)} segments
(begin_padding={self.begin_padding}s, end_padding={self.end_padding}s)...")
91             for idx, (raw_start, raw_end) in enumerate(segments):
92                 # Normalize timestamps to float seconds ? handles float, int,
93                 # plain numeric strings, and colon-separated formats like "MM:SS"
94                 start = self._parse_time(raw_start)
95                 end = self._parse_time(raw_end)
96                 segment_label = f"Segment #{idx+1}/{len(segments)} [{start:.2f}s -
{end:.2f}s]"
97
98                 # Validate the segment times
99                 if start < 0:
100                     logger.warning(f"{segment_label}: start_time is negative
({start:.2f}s). Clamping to 0.")

```

```

101         start = 0.0
102
103     if start >= end:
104         logger.error(f"{segment_label}: start_time ({start:.2f}s) >=
end_time ({end:.2f}s). Skipping invalid segment.")
105         skipped_segments.append((idx, start, end, "start >= end"))
106         continue
107
108     # Check if segment extends beyond video duration
109     if video_duration > 0:
110         if start >= video_duration:
111             logger.error(f"{segment_label}: start_time ({start:.2f}s) is
beyond the video duration ({video_duration:.2f}s). "
112                 f"This segment cannot be extracted ? the video ran
out. Skipping.")
113             skipped_segments.append((idx, start, end, "start beyond video
end"))
114             continue
115
116         if end > video_duration:
117             original_end = end
118             end = video_duration
119             logger.warning(f"{segment_label}: end_time ({original_end:.2f}s)
exceeds video duration ({video_duration:.2f}s). "
120                 f"Clamping end_time to {video_duration:.2f}s. The
segment may be truncated.")
121
122         clip_path = os.path.join(tmp_dir, f"clip_{idx}.mp4")
123
124     # Precise sub-second trim using fast re-encoding to preserve stream
headers
125     padded_start = max(0.0, start - self.begin_padding)
126     padded_end = end + self.end_padding
127
128     # Clamp padded_end to video duration if known
129     if video_duration > 0 and padded_end > video_duration:
130         padded_end = video_duration
131         logger.info(f"{segment_label}: Padded end ({end +
self.end_padding:.2f}s) exceeds video duration. "
132             f"Clamping to {video_duration:.2f}s (end padding
partially applied).")
133
134     trim_duration = padded_end - padded_start
135     logger.info(f"Trimming {segment_label} -> [{padded_start:.3f}s -
{padded_end:.3f}s] (duration: {trim_duration:.2f}s)")
136
137     cmd = [
138         "ffmpeg", "-y",
139         "-ss", str(round(padded_start, 3)),
140         "-to", str(round(padded_end, 3)),
141         "-i", raw_video_path,
142         "-c:v", "libx264",
143         "-preset", "superfast",
144         "-crf", "22",
145         "-c:a", "aac",
146         "-b:a", "128k",
147         clip_path
148     ]
149
150     try:
151         result = subprocess.run(cmd, capture_output=True, check=True)
152         # Verify the output clip was actually created and has content
153         if os.path.exists(clip_path) and os.path.getsize(clip_path) > 0:
154             temp_clips.append(clip_path)
155         else:

```

```

156                 logger.error(f"{segment_label}: FFmpeg exited successfully but
output clip is missing or empty. Skipping.")
157                 skipped_segments.append((idx, start, end, "empty output clip"))
158             except subprocess.CalledProcessError as e:
159                 stderr_msg = e.stderr.decode(errors='replace').strip()
160                 logger.error(f"{segment_label}: FFmpeg trim failed.")
161                 print(f"    Command: {' '.join(cmd)}")
162                 print(f"    FFmpeg stderr: {stderr_msg[-500:]}") # Last 500 chars to
avoid flooding
163                 skipped_segments.append((idx, start, end, f"ffmpeg error:
{stderr_msg[-200:]}"))
164                 continue
165
166             # Report summary of skipped segments
167             if skipped_segments:
168                 logger.info(f"\n[SUMMARY] {len(skipped_segments)} out of {len(segments)}
segments were SKIPPED:")
169                 for seg_idx, seg_start, seg_end, reason in skipped_segments:
170                     logger.info(f"    - Segment #{seg_idx+1} [{seg_start:.2f}s -
{seg_end:.2f}s]: {reason}")
171                 logger.info("")
172
173             if not temp_clips:
174                 raise RuntimeError(
175                     f"[COMPILER ERROR] Failed to generate ANY valid highlight clips. "
176                     f"All {len(segments)} segments were skipped. "
177                     f"Check the errors above for details. "
178                     f"Common causes: video file is corrupted, segments are out of range,
or ffmpeg is not installed."
179                 )
180
181             logger.info(f"Successfully trimmed {len(temp_clips)}/{len(segments)} clips.
Preparing to stitch...")
182
183             # Create a list file for FFmpeg concat demuxer
184             list_file_path = os.path.join(tmp_dir, "clips.txt")
185             with open(list_file_path, "w") as f:
186                 for clip in temp_clips:
187                     # Escape single quotes in path
188                     safe_path = clip.replace("'", "'\\'")
189                     f.write(f"file '{safe_path}'\n")
190
191             logger.info("Stitching clips together...")
192             # Stitch them together without re-encoding (since they were encoded with
identical settings)
193             concat_cmd = [
194                 "ffmpeg", "-y",
195                 "-f", "concat",
196                 "-safe", "0",
197                 "-i", list_file_path,
198                 "-c", "copy",
199                 output_path
200             ]
201
202             try:
203                 result = subprocess.run(concat_cmd, capture_output=True, check=True)
204
205                 # Final verification
206                 if not os.path.exists(output_path) or os.path.getsize(output_path) == 0:
207                     raise RuntimeError(
208                         f"[COMPILER ERROR] FFmpeg concatenation appeared to succeed, "
209                         f"but output file is missing or empty: {output_path}"
210                     )
211
212                 output_size_mb = os.path.getsize(output_path) / (1024 * 1024)

```

```

213         logger.info(f"Dynamic highlight reel saved: {output_path}
({output_size_mb:.1f} MB)")
214         return output_path
215     except subprocess.CalledProcessError as e:
216         stderr_msg = e.stderr.decode(errors='replace').strip()
217         raise RuntimeError(
218             f"[COMPILER ERROR] FFmpeg concatenation failed.\n"
219             f"  Command: {' '.join(concat_cmd)}\n"
220             f"  FFmpeg stderr: {stderr_msg[-500:]}")
221     )
222

```

34. user (2026-06-10T06:04:11.076Z)

```

1      """Audio shuttle-hit ("thwack") detection ? local, dependency-light (numpy + ffmpeg).
2
3      Stages 0?1 of docs/AUDIO_RALLY_DETECTION_PLAN.md: extract the audio track, build a
4      spectral-flux onset envelope (with pre-emphasis to boost the broadband HF transient
5      a racket?shuttle contact makes), threshold it adaptively, and return (timestamp,
6      strength) hit candidates. Pure numpy + stdlib `wave` (NO scipy); ffmpeg (subprocess)
7      only for extraction. Consumed by the E1 go/no-go probe; later feeds rally fusion
8      (ADR-007) if E1 passes. WASB is never modified ? this is a separate signal.
9      """
10     from __future__ import annotations
11
12     import os
13     import wave
14     import tempfile
15     import subprocess
16     from dataclasses import dataclass
17     from typing import List, Tuple
18
19     import numpy as np
20
21
22     @dataclass
23     class HitEvent:
24         t: float          # seconds
25         strength: float    # onset-envelope peak value (arbitrary units)
26
27
28     def extract_audio(video_path: str, out_wav: str = None, sr: int = 16000) -> str:
29         """ffmpeg ? mono 16-bit PCM WAV at `sr`. Returns the wav path (temp if not given).
30         Audio-only decode, so it's fast even for multi-GB video."""
31         if out_wav is None:
32             fd, out_wav = tempfile.mkstemp(suffix=".wav")
33             os.close(fd)
34             cmd = ["ffmpeg", "-y", "-i", video_path, "-vn", "-ac", "1", "-ar", str(sr),
35                  "-c:a", "pcm_s16le", out_wav]
36             subprocess.run(cmd, check=True, stdout=subprocess.DEVNULL,
37                           stderr=subprocess.DEVNULL)
38             return out_wav
39
40     def load_wav(path: str) -> Tuple[np.ndarray, int]:
41         """Read a PCM WAV ? (float32 samples in [-1,1], sample_rate). Mono-downmixes."""
42         with wave.open(path, "rb") as w:
43             sr, n, sw, ch = w.getframerate(), w.getnframes(), w.getsampwidth(),
44             w.getnchannels()
45             raw = w.readframes(n)
46             if sw == 2:
47                 a = np.frombuffer(raw, dtype=np.int16).astype(np.float32) / 32768.0
48             elif sw == 4:
49                 a = np.frombuffer(raw, dtype=np.int32).astype(np.float32) / 2147483648.0
50             elif sw == 1:

```

```

50         a = np.frombuffer(raw, dtype=np.uint8).astype(np.float32) / 128.0 - 1.0
51     else:
52         raise ValueError(f"unsupported sample width {sw}")
53     if ch > 1:
54         a = a.reshape(-1, ch).mean(axis=1)
55     return a, sr
56
57
58     def _preemphasis(x: np.ndarray, coef: float = 0.97) -> np.ndarray:
59         if len(x) < 2:
60             return x.copy()
61         return np.append(x[0], x[1:] - coef * x[:-1])
62
63
64     def onset_envelope(samples: np.ndarray, sr: int, frame: int = 1024,
65                        hop: int = 256, band: Tuple[float, float] = None) ->
66     Tuple[np.ndarray, float]:
67         """Spectral-flux onset strength per hop (pure-numpy STFT). Returns (env,
68         hop_seconds)."""
69         `band=(lo_hz, hi_hz)` restricts the flux to a frequency band ? a shuttle "thwack"
70         is HF-rich while voices/footsteps live lower, so band-limiting (e.g. 3-9 kHz)
71         suppresses non-shuttle transients. Needs sr high enough for hi_hz (Nyquist =
72         sr/2)."""
73         x = _preemphasis(np.asarray(samples, dtype=np.float32))
74         hop_sec = hop / float(sr)
75         if len(x) < frame:
76             return np.zeros(0, dtype=np.float32), hop_sec
77         win = np.hanning(frame).astype(np.float32)
78         n_frames = 1 + (len(x) - frame) // hop
79         idx = np.arange(frame)[None, :] + hop * np.arange(n_frames)[:, None]
80         frames = x[idx] * win
81         mag = np.abs(np.fft.rfft(frames, axis=1))
82         if band is not None:
83             freqs = np.fft.rfftfreq(frame, d=1.0 / sr)
84             keep = (freqs >= band[0]) & (freqs <= band[1])
85             if keep.any():
86                 mag = mag[:, keep] # band-limited flux
87         diff = np.diff(mag, axis=0) # frame-to-frame change
88         flux = np.sqrt((np.maximum(diff, 0.0) ** 2).sum(axis=1)) # positive (onset) flux
89         env = np.concatenate([[0.0], flux]).astype(np.float32) # align to n_frames
90         return env, hop_sec
91
92     def _moving_stats(env: np.ndarray, win: int) -> Tuple[np.ndarray, np.ndarray]:
93         """Sliding-window mean and std (reflect-padded, via cumsum). win forced odd."""
94         if win % 2 == 0:
95             win += 1
96         pad = win // 2
97         p = np.pad(env.astype(np.float64), pad, mode="reflect")
98         c1 = np.cumsum(np.insert(p, 0, 0.0))
99         c2 = np.cumsum(np.insert(p * p, 0, 0.0))
100         s1 = c1[win:] - c1[:-win]
101         s2 = c2[win:] - c2[:-win]
102         mean = s1 / win
103         var = np.maximum(s2 / win - mean * mean, 0.0)
104         return mean[:len(env)], np.sqrt(var)[:len(env)]
105
106     def detect_hits(samples: np.ndarray, sr: int, k: float = 2.5, win_frames: int = 41,
107                    min_gap_sec: float = 0.08, band: Tuple[float, float] = None) ->
108     List[HitEvent]:
109         """Detect impact transients: local maxima of the onset envelope that exceed an
110         adaptive threshold (moving mean + k*std), keeping the strongest within
111         `min_gap_sec`.
```

```

110     `k` is the sensitivity knob (lower = more hits); `band` band-limits the onset flux
111     to the shuttle-click frequencies (e.g. (3000, 9000))."""
112     env, hop_sec = onset_envelope(samples, sr, band=band)
113     if len(env) < 3:
114         return []
115     mean, std = _moving_stats(env, win_frames)
116     thr = mean + k * std
117     is_peak = (env[1:-1] >= env[:-2]) & (env[1:-1] > env[2:]) & (env[1:-1] > thr[1:-1])
118     peak_idx = np.nonzero(is_peak)[0] + 1
119     if peak_idx.size == 0:
120         return []
121     # enforce min inter-peak gap: greedily keep strongest first
122     min_gap = max(1, int(round(min_gap_sec / hop_sec)))
123     order = peak_idx[np.argsort(-env[peak_idx])]
124     taken = np.zeros(len(env), dtype=bool)
125     kept = []
126     for i in order:
127         lo, hi = max(0, i - min_gap), min(len(env), i + min_gap + 1)
128         if taken[lo:hi].any():
129             continue
130         taken[i] = True
131         kept.append(i)
132     kept.sort()
133     return [HitEvent(t=float(i * hop_sec), strength=float(env[i])) for i in kept]
134
135
136 def detect_hits_in_video(video_path: str, sr: int = 16000, **kw) -> List[HitEvent]:
137     """Convenience: extract audio from a video and detect hits. Cleans up the temp
138     wav."""
139     wav = extract_audio(video_path, sr=sr)
140     try:
141         samples, real_sr = load_wav(wav)
142         return detect_hits(samples, real_sr, **kw)
143     finally:
144         try:
145             os.remove(wav)
146         except OSError:
147             pass

```

35. user (2026-06-10T06:04:29.973Z)

```

1     import os
2     import time
3     import json
4     import logging
5     from typing import Tuple
6     from google import genai
7     from google.genai import types
8     from backend.providers.base import AIVideoProvider
9
10    logger = logging.getLogger(__name__)
11
12    # The Gemini File API caps uploads at 2GB; stay well under it. Callers must chunk so
13    # each uploaded clip is below this. Oversize uploads are refused (not silently
14    # attempted).
15    MAX_UPLOAD_MB = 500
16
17    class GeminiProvider(AIVideoProvider):
18        def __init__(self, api_key: str, model_name: str):
19            super().__init__(api_key, model_name)
20            # Initialize client lazily to avoid connection/auth issues during class
21            # instantiation
22            self._client = None

```



```

22     @property
23     def client(self):
24         if self._client is None:
25             self._client = genai.Client(api_key=self.api_key)
26         return self._client
27
28     def analyze_video(self, video_path: str, prompt: str,
29                       chunk_duration: float = 0.0,
30                       label: str = "") -> Tuple[list, dict]:
31         log_prefix = f"Gemini {label}" if label else "Gemini"
32
33         metrics = {
34             "upload_time": 0.0,
35             "process_time": 0.0,
36             "gen_time": 0.0
37         }
38
39         # 0. Refuse oversize uploads (Gemini File API limit is 2GB; we cap lower for
safety/speed)
40         try:
41             size_mb = os.path.getsize(video_path) / (1024 * 1024)
42         except OSError:
43             size_mb = 0.0
44         if size_mb > MAX_UPLOAD_MB:
45             logger.error(f"[{log_prefix}] {video_path} is {size_mb:.0f}MB >
{MAX_UPLOAD_MB}MB cap ? ")
46             return f"chunk smaller. Skipping upload."
47         return [], metrics
48
49         # 1. Upload (retry transient network failures, e.g. WinError 10053 / reset
connections)
50         logger.info(f"[{log_prefix}] Uploading {video_path}...")
51         t_upload_start = time.time()
52         video_file = None
53         for attempt in range(3):
54             try:
55                 video_file = self.client.files.upload(file=video_path)
56                 metrics["upload_time"] = time.time() - t_upload_start
57                 logger.info(f"[{log_prefix}] Upload complete in
{metrics['upload_time']:.1f}s. URI: {video_file.uri}")
58                 break
59             except Exception as e:
60                 logger.warning(f"[{log_prefix}] Upload attempt {attempt + 1}/3 failed:
{e}")
61                 if attempt < 2:
62                     time.sleep(2 * (attempt + 1))
63         if video_file is None:
64             logger.error(f"[{log_prefix}] Upload failed after 3 attempts.")
65             return [], metrics
66
67         # 2. Wait for processing on Google servers
68         t_process_start = time.time()
69         try:
70             while video_file.state.name == "PROCESSING":
71                 logger.info(f"[{log_prefix}] Processing on Google servers. Elapsed:
{time.time()-t_process_start:.1f}s. Retrying in 10s...")
72                 time.sleep(10)
73                 video_file = self.client.files.get(name=video_file.name)
74                 if video_file.state.name == "FAILED":
75                     logger.error(f"[{log_prefix}] Video processing failed:
{video_file.error.message}")
76                 try:
77                     self.client.files.delete(name=video_file.name)
78                 except Exception:
79                     pass

```

```

80         return [], metrics
81     except Exception as e:
82         logger.error(f"[{log_prefix}] Status check failed: {e}")
83         try:
84             self.client.files.delete(name=video_file.name)
85         except Exception:
86             pass
87         return [], metrics
88
89         metrics["process_time"] = time.time() - t_process_start
90         logger.info(f"[{log_prefix}] Processing finished in
{metrics['process_time']:.1f}s.")
91
92         # 3. Call Gemini (retry transient 503/UNAVAILABLE model-overload + network
blips)
93         logger.info(f"[{log_prefix}] Submitting multimodal request to
{self.model_name}...")
94         t_gen_start = time.time()
95         raw_text = None
96         for attempt in range(3):
97             try:
98                 response = self.client.models.generate_content(
99                     model=self.model_name,
100                     contents=[video_file, prompt],
101                     config=types.GenerateContentConfig(
102                         response_mime_type="application/json",
103                         temperature=0.0
104                     )
105                 )
106                 raw_text = response.text
107                 metrics["gen_time"] = time.time() - t_gen_start
108                 logger.info(f"[{log_prefix}] Response received in
{metrics['gen_time']:.1f}s:\n{raw_text}")
109                 break
110             except Exception as e:
111                 logger.warning(f"[{log_prefix}] generate attempt {attempt + 1}/3 failed:
{e}")
112                 if attempt < 2:
113                     time.sleep(3 * (attempt + 1))
114             try:
115                 self.client.files.delete(name=video_file.name)
116         except Exception:
117             pass
118
119         if not raw_text:
120             return [], metrics
121
122         # 4. Parse JSON
123         try:
124             segments = json.loads(raw_text)
125             if not isinstance(segments, list):
126                 logger.error(f"[{log_prefix}] Response is not a JSON array.")
127             return [], metrics
128         except Exception as e:
129             logger.error(f"[{log_prefix}] JSON parse failed: {e}")
130             return [], metrics
131
132

```

36. user (2026-06-10T06:04:31.289Z)

```

run_process_and_stitch.py:87: end_padding = cfg.stitching.end_padding
run_process_and_stitch.py:88: begin_padding = cfg.stitching.begin_padding
run_process_and_stitch.py:89: min_shot_count = cfg.stitching.min_shot_count
run_process_and_stitch.py:94: bpad_input = input(f"Enter begin_padding in seconds

```

```

(default {begin_padding}): ").strip()
run_process_and_stitch.py:96:         begin_padding = float(bpad_input)
run_process_and_stitch.py:98:         pad_input = input(f"Enter end_padding in seconds
(default {end_padding}): ").strip()
run_process_and_stitch.py:100:         end_padding = float(pad_input)
run_process_and_stitch.py:102:         shot_input = input(f"Enter min_shot_count to filter by
(default {min_shot_count}): ").strip()
run_process_and_stitch.py:104:         min_shot_count = int(shot_input)
run_process_and_stitch.py:108:     print(f"\nCompiling highlights (with
begin_padding={begin_padding}s, end_padding={end_padding}s,
min_shot_count={min_shot_count})...", flush=True)
run_process_and_stitch.py:110:     # Filter segments where shot_count >= min_shot_count
run_process_and_stitch.py:121:         if shot_count >= min_shot_count:
run_process_and_stitch.py:127:         print(f"[FILTER] Excluded {skipped_count} play segments
with fewer than {min_shot_count} shots.", flush=True)
run_process_and_stitch.py:134:     compiler = VideoCompiler(output_dir, end_padding=end_padding,
begin_padding=begin_padding)
backend\main.py:201:     compiler = VideoCompiler(output_dir)
backend\main.py:324:     compiler = VideoCompiler(out_dir)
backend\config\models.py:86:     begin_padding: float = 0.5
backend\config\models.py:87:     end_padding: float = 0.5
backend\config\models.py:88:     use_cache: bool = False
backend\config\models.py:89:     min_shot_count: int = 1
tests\test_config.py:131:     assert cfg.get("begin_padding") == 0.5        # lives inside
`stitching`
backend\pipeline\stitcher\compiler.py:10:     def __init__(self, output_dir: str, end_padding:
float = 1.0, begin_padding: float = 0.5):
backend\pipeline\stitcher\compiler.py:12:         self.end_padding = end_padding
backend\pipeline\stitcher\compiler.py:13:         self.begin_padding = begin_padding
backend\pipeline\stitcher\compiler.py:67:         into a seamless highlights file, applying
end_padding to prevent abrupt endings.
backend\pipeline\stitcher\compiler.py:90:         logger.info(f"Trimming {len(segments)}
segments (begin_padding={self.begin_padding}s, end_padding={self.end_padding}s)...")
backend\pipeline\stitcher\compiler.py:125:         padded_start = max(0.0, start -
self.begin_padding)
backend\pipeline\stitcher\compiler.py:126:         padded_end = end + self.end_padding
backend\pipeline\stitcher\compiler.py:131:         logger.info(f"{segment_label}:
Padded end ({end + self.end_padding:.2f}s) exceeds video duration. ")
backend\pipeline\segmenters\gemini.py:380:         begin_pad = stitch_cfg.get("begin_padding",
0.5)
backend\pipeline\segmenters\gemini.py:381:         end_pad = stitch_cfg.get("end_padding", 0.5)

```

37. user (2026-06-10T06:04:43.023Z)

```

1     from typing import Dict, Type
2     from backend.providers.base import AIVideoProvider
3     from backend.providers.gemini import GeminiProvider
4
5     class ProviderRegistry:
6         def __init__(self):
7             self._registry: Dict[str, Type[AIVideoProvider]] = {}
8             # Register Gemini provider
9             self.register("gemini", GeminiProvider)
10
11         def register(self, name: str, provider_cls: Type[AIVideoProvider]):
12             self._registry[name] = provider_cls
13
14         def get(self, name: str, api_key: str, model_name: str) -> AIVideoProvider:
15             if name not in self._registry:
16                 raise KeyError(f"AI Provider '{name}' is not registered in the
ProviderRegistry. Available: {list(self._registry.keys())}")
17             return self._registry[name](api_key=api_key, model_name=model_name)
18
19         def list_available(self) -> Dict[str, str]:
20             return {name: cls.__doc__ or f"{name.capitalize()} provider" for name, cls in

```

```

self._registry.items()}
21
22     provider_registry = ProviderRegistry()
23

```

38. user (2026-06-10T06:04:43.980Z)

```

1     from abc import ABC, abstractmethod
2     from typing import Tuple
3
4     class AIVideoProvider(ABC):
5         def __init__(self, api_key: str, model_name: str):
6             self.api_key = api_key
7             self.model_name = model_name
8
9         @abstractmethod
10        def analyze_video(self, video_path: str, prompt: str,
11                          chunk_duration: float = 0.0,
12                          label: str = "") -> Tuple[list, dict]:
13            """Sends a video file to the AI model along with a prompt, and returns a tuple
14            of
15            (segments_list, metrics_dict).
16
17            Args:
18                video_path: Path to the video file to analyze.
19                prompt: Prompt instructions.
20                chunk_duration: Duration of this chunk in seconds, or 0.0 if not chunked.
21                label: Diagnostic label for logging.
22
23            Returns:
24                Tuple containing:
25                - A list of detected segment dicts (e.g. [{"start_time": ..., "end_time":
26                ...}])
27                - A dictionary of performance metrics (e.g. {"upload_time": ...,
28                "process_time": ..., "gen_time": ...})
29            """
30            pass
31

```

39. user (2026-06-10T06:04:44.460Z)

```

1     from typing import Dict, Type
2     from backend.sports.base import SportProfile
3     from backend.sports.badminton import BadmintonSportProfile
4
5     class SportRegistry:
6         def __init__(self):
7             self._registry: Dict[str, Type[SportProfile]] = {}
8             # Automatically register Badminton
9             self.register(BadmintonSportProfile)
10
11        def register(self, profile_cls: Type[SportProfile]):
12            inst = profile_cls()
13            self._registry[inst.name] = profile_cls
14
15        def get(self, name: str) -> SportProfile:
16            if name not in self._registry:
17                raise KeyError(f"Sport '{name}' is not registered in the SportRegistry.
18                Available: {list(self._registry.keys())}")
19            return self._registry[name]()
20
21        def list_available(self) -> Dict[str, str]:
22            return {name: cls.__doc__ or f"{name.capitalize()} sport profile" for name, cls
23            in self._registry.items()}
24

```

```
23     sport_registry = SportRegistry()
24
```

40. user (2026-06-10T06:04:45.116Z)

```
1     from backend.sports.base import SportProfile
2
3     class BadmintonSportProfile(SportProfile):
4         @property
5         def name(self) -> str:
6             return "badminton"
7
8         def get_prompt(self, chunk_duration: float = 0.0) -> str:
9             base_prompt = ""
10             You are a highly precise computer vision data extraction engine specializing in
11             badminton match analysis. Your exclusive objective is to parse a full-length match video and
12             generate a flawless JSON array containing exact rally boundaries and shot metrics for automated
13             highlight generation.
14
15             Core Directives
16             Detect every single active rally in the video. Do not miss any live gameplay.
17             Extract high-precision timestamps as TOTAL DECIMAL SECONDS from the start of the video
18             for exact frame-accurate cuts.
19             Estimate the number of shots exchanged (net crossings) and provide a confidence score
20             for that count, acknowledging frame-sampling limitations.
21
22             CRITICAL TIMESTAMP FORMAT RULES:
23             - All start_time and end_time values MUST be expressed as TOTAL SECONDS from the
24             beginning of the video, as a decimal number.
25             - Example: A rally starting at 1 minute 14 seconds into the video MUST be written as
26             74.0, NOT as "1:14" or 114.
27             - Example: A rally starting at 2 minutes 30.5 seconds MUST be written as 150.5, NOT as
28             "2:30.5" or 230.5.
29             - NEVER use MM:SS or HH:MM:SS colon-separated format. NEVER concatenate minutes and
30             seconds into a single number without converting.
31             - The downstream pipeline will BREAK if timestamps are not in total decimal seconds.
32
33             Positive Constraints: Rally Boundaries
34             Start Condition: The exact fractional second the server's racket makes contact with the
35             shuttlecock.
36             End Condition: The rally concludes ONLY when both of the following physical criteria are
37             met sequentially: First, the shuttlecock hits the court surface, strikes the net and drops, or
38             lands visibly out of bounds. Second, a 0.3 to 0.5-second buffer has elapsed where both players
39             have visibly ceased all kinetic momentum (no lunging, sliding, or racket follow-through).
40
41             Negative Constraints: Anomaly Rejection (CRITICAL)
42             Replay Rejection: You must strictly ignore and filter out all slow-motion instant
43             replays. Replays are not live rallies.
44             Dead-Time Rejection: Ignore all Hawkeye/line-call challenges, interval coaching, towel
45             breaks, court mopping, changing shuttles, and player celebrations.
46             Pre-serve Rituals: Do not start the timestamp during the players' setup or stance. Wait
47             for physical racket-to-shuttle contact.
48
49             Shot Tracking & Data Confidence
50             Shot Definition: Count every successful traverse of the shuttlecock across the net
51             (serves, clears, smashes, drops, drives). Do not count the final grounding or pre-net faults.
52             Confidence Scoring: Because the shuttlecock moves at extreme velocities that may fall
53             between sampled frames, provide a shot_count_confidence metric (Low, Medium, High). Base your
54             estimate on player biomechanics, positioning, and racket swings if the shuttle itself is
55             blurred.
56
57             Strict Output Schema
58             Output strictly in the following JSON array format. Do not include markdown formatting,
59             introductory text, conversational filler, or explanations. Any deviation from this schema will
60             break the downstream pipeline.
```

```

39     Remember: start_time and end_time are TOTAL SECONDS (decimal), not MM:SS.
40
41     [
42     {
43         "start_time": 74.0,
44         "end_time": 81.3,
45         "shuttle_exchange_count": 11,
46         "shot_count_confidence": "High"
47     },
48     {
49         "start_time": 150.5,
50         "end_time": 174.2,
51         "shuttle_exchange_count": 24,
52         "shot_count_confidence": "Medium"
53     }
54     ]
55     """
56
57     if chunk_duration > 0.0:
58         duration_context = (
59             f"IMPORTANT VIDEO CONTEXT: This video segment is exactly
{chunk_duration:.1f} seconds long. "
60             f"All timestamps MUST be between 0 and {chunk_duration:.1f}. "
61             f"Any timestamp beyond {chunk_duration:.1f}s is INVALID.\n\n"
62         )
63         return duration_context + base_prompt
64     return base_prompt
65
66     def get_relevance_config(self) -> dict:
67         return {
68             "court_green_lower": [35, 40, 40],
69             "court_green_upper": [85, 255, 255],
70             "court_pixels_min_ratio": 0.05
71         }

```

41. user (2026-06-10T06:04:45.522Z)

```

1     from abc import ABC, abstractmethod
2
3     class SportProfile(ABC):
4         @property
5         @abstractmethod
6         def name(self) -> str:
7             """The identifier of the sport (e.g., 'badminton')."""
8             pass
9
10        @abstractmethod
11        def get_prompt(self, chunk_duration: float = 0.0) -> str:
12            """Full system prompt for the AI model, including output schema."""
13            pass
14
15        @abstractmethod
16        def get_relevance_config(self) -> dict:
17            """Sport-specific visual validation params (court colors, thresholds, etc.)."""
18            pass
19

```

42. user (2026-06-10T06:04:56.922Z)

```

1     """Per-video observability reports for the indexer.
2
3     `obsreport` is a SOFT dependency: if it is not installed, this module degrades to a
4     no-op and the pipeline runs exactly as before. The single entrypoint
5     `emit_indexer_report(...)` is designed to be called from a `finally` block and
6     NEVER raises ? emitting a report must not be able to break a processing run.

```

```

7     """
8     from __future__ import annotations
9
10    import os
11    from typing import Any, Dict, List, Optional
12
13    try: # soft dependency ? never hard-fail the pipeline on a reporting import
14        import obsreport # noqa: F401
15        from . import harvest
16
17        OBSREPORT_AVAILABLE = True
18    except Exception: # pragma: no cover - exercised only when obsreport is absent
19        OBSREPORT_AVAILABLE = False
20
21
22    def report_enabled(config: Optional[Dict[str, Any]]) -> bool:
23        return bool((config or {}).get("report", {}).get("enabled", True))
24
25
26    def emit_indexer_report(
27        *,
28        db,
29        video_id: str,
30        video_path: Optional[str],
31        config: Optional[Dict[str, Any]],
32        val_results: List[Any],
33        val_context: Optional[Dict[str, Any]] = None,
34        segmenter_requested: Optional[str] = None,
35        segmenter_actual: Optional[str] = None,
36        was_reingest: bool = False,
37        segmentation_ran: bool = True,
38    ) -> Optional[Dict[str, str]]:
39        """Harvest + write the report. Returns the written file paths, or None if
40        reporting is unavailable/disabled/failed. Never raises."""
41        config = config or {}
42        # Accept either a plain dict or a typed config object (e.g.
43        backend.config.AppConfig).
44        # Normalize to a plain nested dict so harvest's dict access works uniformly.
45        if hasattr(config, "model_dump"):
46            try:
47                config = config.model_dump(mode="json")
48            except Exception:
49                config = {}
50        if not OBSREPORT_AVAILABLE or not report_enabled(config):
51            return None
52        try:
53            play_segments = db.query_play_segments(video_id, {}) if db is not None else []
54            candidates = db.query_candidate_segments(video_id) if db is not None else []
55            rec = harvest.record_run(
56                video_id=video_id,
57                video_path=video_path,
58                config=config,
59                val_results=val_results,
60                val_context=val_context,
61                play_segments=play_segments,
62                candidates=candidates,
63                segmenter_requested=segmenter_requested,
64                segmenter_actual=segmenter_actual,
65                was_reingest=was_reingest,
66                segmentation_ran=segmentation_ran,
67            )
68            output_dir = (config.get("report") or {}).get("output_dir")
69            res = rec.write(output_dir=output_dir)
70            return res.paths or None
71        except Exception: # pragma: no cover - defensive; reporting must never break a run

```

```
71         return None
72
```

43. user (2026-06-10T06:04:57.230Z)

```
1     """Build an observability report for one indexer run.
2
3     Pure-ish harvest: given the validation results/context, the persisted segments,
4     and run metadata, populate an `obsreport.RunRecorder` with the raw signals. The
5     *footgun rules* (which turn signals into flags) live declaratively in spec.yaml,
6     so this module only records facts faithfully ? it does not decide warnings.
7
8     Imported only after `backend.reporting` confirms obsreport is installed, so it may
9     import obsreport freely.
10    """
11    from __future__ import annotations
12
13    import os
14    from pathlib import Path
15    from typing import Any, Dict, List, Optional
16
17    try:
18        import obsreport # soft dep; only needed for RunRecorder / load_spec at runtime
19    except ImportError:
20        obsreport = None # type: ignore[assignment]
21
22    _AI_SEGMENTERS = {"yolo_hybrid", "tracknet_hybrid", "wasb_hybrid", "gemini"}
23    _SPEC_PATH = Path(__file__).resolve().parent / "spec.yaml"
24
25    def load_spec() -> "obsreport.ReportSpec":
26        return obsreport.load_spec(_SPEC_PATH)
27
28
29    # --- media metadata -----
30    def video_meta_from_context(val_context: Optional[Dict[str, Any]], video_path:
Optional[str]) -> Dict[str, Any]:
31        """Derive VideoMeta fields from the validators' shared ffprobe_meta. fps_source
32        distinguishes a real probe from the absence of one (the fallback footgun)."""
33        meta: Dict[str, Any] = {"fps_source": "unknown"}
34        if video_path and os.path.exists(video_path):
35            try:
36                meta["size_mb"] = round(os.path.getsize(video_path) / (1024 * 1024), 1)
37            except OSError:
38                pass
39
40
41        ffmata = (val_context or {}).get("ffprobe_meta") or {}
42        streams = ffmata.get("streams") or []
43        fmt = ffmata.get("format") or {}
44        if streams:
45            s = streams[0]
46            w, h = s.get("width"), s.get("height")
47            meta["width"], meta["height"] = w, h
48            if w and h:
49                meta["resolution"] = f"{w}x{h}"
50            fps_str = s.get("r_frame_rate", "0/1")
51            try:
52                num, den = (int(x) for x in fps_str.split("/"))
53                if den:
54                    meta["fps"] = round(num / den, 2)
55                    meta["fps_source"] = "probed"
56            except (ValueError, ZeroDivisionError):
57                pass
58            if s.get("codec_name"):
59                meta["container"] = fmt.get("format_name") or s.get("codec_name")
```



```

60     if fmt.get("duration"):
61         try:
62             meta["duration_s"] = round(float(fmt["duration"]), 1)
63         except (ValueError, TypeError):
64             pass
65     if fmt.get("format_name"):
66         meta["container"] = fmt.get("format_name")
67
68     # Audio readiness signal (for Input Quality UX extension per review item 2).
69     # Guidelines make "record with audio on, mic unobstructed" a first-class requirement
70     # (ADR-007) because shuttle thwacks are a strong rally cue for recall.
71     # We only record presence here (cheap, from ffmpeg already in context); the actual
72     # thwack onset probe (pipeline/audio/hit_detector) remains a diagnostic /
calibration
73     # tool for now.
74     audio_streams = [s for s in streams if (s or {}).get("codec_type") == "audio"]
75     meta["audio_present"] = len(audio_streams) > 0
76     return meta
77
78
79     # --- stage builders -----
80     def _validation_stage(rec: "obsreport.RunRecorder", val_results: List[Any]) -> None:
81         with rec.stage("validation") as st:
82             total = len(val_results)
83             passed = sum(1 for r in val_results if getattr(r, "passed", False))
84             st.add_metric("checks_passed", passed)
85             st.add_metric("checks_total", total)
86             failures = [r for r in val_results if not getattr(r, "passed", True)]
87             for r in failures:
88                 st.messages.append(f"FAILED {getattr(r, 'validator_name', '?')}: {getattr(r,
'message', '')}")
89             st.status = "error" if failures else ("ok" if total else "not_run")
90
91
92     def _segmentation_stage(
93         rec: "obsreport.RunRecorder",
94         play_segments: List[Dict[str, Any]],
95         candidates: List[Dict[str, Any]],
96         segmenter_requested: Optional[str],
97         segmenter_actual: Optional[str],
98         config_default: Optional[str],
99         was_reingest: bool,
100         ran: bool,
101     ) -> None:
102         with rec.stage("segmentation") as st:
103             if not ran:
104                 st.status = "not_run"
105             seg_count = len(play_segments)
106             total_play = round(sum(float(s.get("duration", 0.0)) for s in play_segments), 1)
107             detector_tag = None
108             if play_segments:
109                 detector_tag = (play_segments[0].get("metadata") or {}).get("detector")
110             st.add_metric("candidate_windows", len(candidates))
111             st.add_metric("segment_count", seg_count, audience="both")
112             st.add_metric("total_play_s", total_play, unit="s")
113             # "diverges" only when a segmenter actually RAN, the user did NOT explicitly
114             # pick it, and it differs from config's default (the genuine surprise case).
115             # An explicit --segmenter override or a halted run is not a divergence.
116             matches_default = True
117             if ran and segmenter_actual and config_default:
118                 matches_default = (segmenter_actual == config_default)
119             st.details.update({
120                 "segmenter_requested": segmenter_requested,
121                 "segmenter_actual": segmenter_actual,
122                 "config_default": config_default,

```

```

123         "segmenter_explicitly_requested": segmenter_requested is not None,
124         "matches_config_default": matches_default,
125         "was_reingest": was_reingest,
126         "detector": detector_tag,
127     })
128     # Surface metadata trustworthiness at the data level (not just in a warning):
129     # the motion baseline fabricates per-segment metadata; everything else measures
130     it.
131     if ran and segmenter_actual:
132         st.details["metadata_source"] = (
133             "synthetic/heuristic ? NOT measured" if segmenter_actual == "motion"
134             else "measured (detector/AI)"
135         )
136
137     def _ai_stage(rec: "obsreport.RunRecorder", config: Dict[str, Any], segmenter_actual:
Optional[str],
138                 play_segments: List[Dict[str, Any]], ran: bool) -> None:
139         """Only meaningful for AI-based segmenters. Records whether the provider key was
140         present ? the signal the silent_provider_noop rule keys on."""
141         ai_based = segmenter_actual in _AI_SEGMENTERS
142         provider = (config.get("ai_provider") or "gemini")
143         key_env = f"{provider.upper()}_API_KEY"
144         with rec.stage("ai_handoff") as st:
145             if not ran or not ai_based:
146                 st.status = "not_run" if not ran else "skipped"
147             st.details.update({
148                 "ai_based": ai_based,
149                 "provider": provider if ai_based else None,
150                 "model": config.get("ai_model") if ai_based else None,
151                 "api_key_present": bool(os.environ.get(key_env)) if ai_based else None,
152             })
153             if ai_based:
154                 st.add_metric("confirmed_segments", len(play_segments))
155
156
157     def _highlights(rec: "obsreport.RunRecorder", play_segments: List[Dict[str, Any]],
158                   video_meta: Dict[str, Any]) -> None:
159         total = round(sum(float(s.get("duration", 0.0)) for s in play_segments), 1)
160         coverage = None
161         dur = video_meta.get("duration_s")
162         if dur:
163             coverage = round(100.0 * total / dur, 1)
164         notes = []
165         if not play_segments:
166             notes.append("No highlights were produced for this video.")
167         rec.set_highlights(segment_count=len(play_segments), total_highlight_s=total,
168                           coverage_pct=coverage, notes=notes)
169
170
171     # --- top-level -----
172     def record_run(
173         *,
174         video_id: str,
175         video_path: Optional[str],
176         config: Dict[str, Any],
177         val_results: List[Any],
178         val_context: Optional[Dict[str, Any]],
179         play_segments: List[Dict[str, Any]],
180         candidates: List[Dict[str, Any]],
181         segmenter_requested: Optional[str] = None,
182         segmenter_actual: Optional[str] = None,
183         was_reingest: bool = False,
184         segmentation_ran: bool = True,
185         tool_version: Optional[str] = None,

```

```

186         spec: Optional["obsreport.ReportSpec"] = None,
187         **recorder_kwargs: Any,
188     ) -> "obsreport.RunRecorder":
189         """Assemble a finalized recorder for an indexer run (does NOT write)."""
190         rec = obsreport.RunRecorder(
191             "badminton-highlight-indexer",
192             tool_version=tool_version,
193             video_id=video_id,
194             video_id_kind="file_md5",
195             video_path=video_path,
196             spec=spec if spec is not None else load_spec(),
197             **recorder_kwargs,
198         )
199         vm = video_meta_from_context(val_context, video_path)
200         # If we never probed an fps but segmentation still ran, the segmenter must have
201         # used its internal fallback (30fps / 1280px) ? surface that as the footgun it is,
202         # AND show the assumed numeric value so the report reads "30 (fallback)" not "?
203         # (fallback)".
204         if vm.get("fps_source") == "unknown" and segmentation_ran:
205             vm["fps_source"] = "fallback"
206             if vm.get("fps") is None:
207                 vm["fps"] = 30.0 # matches the runners' fallback (_probe_fps_width -> 30.0
208                 / 1280)
209         rec.set_video_meta(**vm)
210         _validation_stage(rec, val_results)
211         config_default = (config.get("indexing") or {}).get("default_segmenter")
212         _segmentation_stage(rec, play_segments, candidates, segmenter_requested,
213                             segmenter_actual, config_default, was_reingest,
214                             segmentation_ran)
215         _ai_stage(rec, config, segmenter_actual, play_segments, segmentation_ran)
216         _highlights(rec, play_segments, vm)
217         rec.finalize()
218         return rec

```

44. user (2026-06-10T06:05:05.433Z)

```

1     """Annotation provider registry ? pluggable sources of ground-truth rally labels.
2
3     An `AnnotationProvider` turns a video reference into a list of `(start, end)` rally
4     intervals ? the **same shape** that `backend.eval.annotations.load_annotations`
5     produces and `backend.eval.metrics` consumes. That symmetry is the whole point:
6     golden-set ground truth can come from a CSV file, a human annotation vendor, or an
7     automated oracle behind one interface, so the regression flow
8     (`backend.eval.regression`) never needs to know where labels came from.
9
10    Mirrors the existing registry pattern in `backend/providers/base.py` and
11    `backend/indexer/base.py`. See `docs/GOLDEN_SET_VENDURING.md` (and ADR-006) for the
12    design rationale.
13    """
14    from abc import ABC, abstractmethod
15    from typing import Dict, List, Optional, Tuple, Type
16
17    # An (start, end) interval in seconds ? identical to backend.eval.metrics.Interval.
18    Interval = Tuple[float, float]
19
20    # Job lifecycle statuses returned by poll().
21    QUEUED = "queued"
22    RUNNING = "running"
23    DONE = "done"
24    FAILED = "failed"
25    STATUSES = (QUEUED, RUNNING, DONE, FAILED)
26

```

```

27
28 class AnnotationProvider(ABC):
29     """Abstract source of ground-truth rally intervals for a video.
30
31     The three-method contract (submit ? poll ? fetch) accommodates both
32     synchronous providers (a file already on disk) and asynchronous ones (a human
33     vendor that takes days). Synchronous providers can use the `annotate()``
34     convenience, which runs the full cycle in one call.
35     """
36
37     @property
38     @abstractmethod
39     def name(self) -> str:
40         """Unique registry key for this provider (e.g. 'file', 'vendor', 'auto')."""
41
42     @abstractmethod
43     def submit(self, video_ref: str, guideline: Optional[str] = None,
44               sport: str = "badminton") -> str:
45         """Submit a video for annotation; return an opaque job_id.
46
47         `video_ref` is provider-defined (a file path, a video_id/MD5, a URL).
48         `guideline` is the annotation spec (what counts as a rally start/end).
49         """
50
51     @abstractmethod
52     def poll(self, job_id: str) -> str:
53         """Return one of STATUSES for the given job."""
54
55     @abstractmethod
56     def fetch(self, job_id: str) -> List[Interval]:
57         """Return the completed annotations as sorted `(start, end)` intervals.
58
59         Should raise if the job is not DONE / the labels are unavailable, rather
60         than returning an empty list (which would silently skew downstream scores).
61         """
62
63     def annotate(self, video_ref: str, guideline: Optional[str] = None,
64               sport: str = "badminton") -> List[Interval]:
65         """Convenience: submit + fetch for a provider that completes synchronously.
66
67         Async providers (e.g. a human vendor with a multi-day SLA) should be driven
68         via submit/poll/fetch explicitly instead ? calling this on one will raise
69         once the job isn't immediately DONE.
70         """
71         job_id = self.submit(video_ref, guideline=guideline, sport=sport)
72         status = self.poll(job_id)
73         if status != DONE:
74             raise RuntimeError(
75                 f"{self.name}.annotate(): job '{job_id}' is '{status}', not DONE. "
76                 f"Use submit/poll/fetch for asynchronous providers."
77             )
78         return self.fetch(job_id)
79
80
81 class AnnotationProviderRegistry:
82     """Registry of annotation providers, keyed by each provider's `name`."""
83
84     def __init__(self):
85         self._providers: Dict[str, Type[AnnotationProvider]] = {}
86
87     def register(self, provider_cls: Type[AnnotationProvider]) ->
Type[AnnotationProvider]:
88         """Register a provider class under its `name`. Returns the class (decorator-
friendly)."""
89         try:

```

```

90         key = provider_cls().name
91     except Exception:
92         # Providers needing constructor args fall back to the class name.
93         key = getattr(provider_cls, "__name__", "unknown")
94     self._providers[key] = provider_cls
95     return provider_cls
96
97     def get(self, name: str, **kwargs) -> AnnotationProvider:
98         """Instantiate the named provider, passing through any constructor kwargs."""
99         cls = self._providers.get(name)
100         if cls is None:
101             raise KeyError(
102                 f"Unknown annotation provider '{name}'. Available:
{self.list_available()}"
103             )
104         return cls(**kwargs)
105
106     def list_available(self) -> List[str]:
107         return list(self._providers.keys())
108
109
110 annotation_registry = AnnotationProviderRegistry()
111

```

45. user (2026-06-10T06:05:05.905Z)

```

1     """FileProvider ? ground-truth rally labels from pre-made CSV files.
2
3     This formalizes today's manual / ShuttleSet annotation flow behind the
4     `AnnotationProvider` interface: the "annotation job" is simply a CSV that already
5     exists on disk (the same `start,end` format `run_eval.py --scaffold` writes). The
6     job_id IS the resolved CSV path, so submit/poll/fetch all operate on that path.
7
8     It's the synchronous, zero-cost provider used as the default golden-set source and
9     as the stand-in in tests/CI before a human-vendor or automated provider is wired in.
10    """
11    import os
12    from typing import List, Optional
13
14    from backend.annotations.base import (
15        AnnotationProvider,
16        annotation_registry,
17        Interval,
18        DONE,
19        FAILED,
20    )
21    from backend.eval.annotations import load_annotations
22
23
24    class FileProvider(AnnotationProvider):
25        def __init__(self, csv_dir: Optional[str] = None, ext: str = ".csv"):
26            """
27            Args:
28                csv_dir: if set, a bare `video_ref` (not itself a path) is resolved to
29                `

```

```

40         """Map a video_ref to a CSV path.
41
42         An existing path is used as-is; otherwise it's treated as a video_id and
43         looked up under csv_dir (if configured).
44         """
45         if os.path.isfile(video_ref):
46             return video_ref
47         if self.csv_dir:
48             return os.path.join(self.csv_dir, f"{video_ref}{self.ext}")
49         return video_ref
50
51     def submit(self, video_ref: str, guideline: Optional[str] = None,
52               sport: str = "badminton") -> str:
53         # Nothing to do ? the label file is expected to already exist. The job_id
54         # is the resolved CSV path. (guideline/sport are irrelevant for a file.)
55         return self._resolve(video_ref)
56
57     def poll(self, job_id: str) -> str:
58         return DONE if os.path.isfile(job_id) else FAILED
59
60     def fetch(self, job_id: str) -> List[Interval]:
61         if not os.path.isfile(job_id):
62             raise FileNotFoundError(f"annotation file not found: {job_id}")
63         return load_annotations(job_id)
64
65
66     annotation_registry.register(FileProvider)
67

```

46. user (2026-06-10T06:05:06.252Z)

```

1     """AutomatedProvider ? the fast-tier "oracle" annotation source (GS-4, interface stub).
2
3     The fast/CI tier of the golden-set flow needs an *automated* labeler that turns a
4     video into rally `(start, end)` intervals quickly and cheaply, as a smoke alarm
5     between authoritative human-vendor runs (the slow tier).
6
7     This slice ships the **interface only**, deliberately model-agnostic: an
8     `AutomatedProvider` that delegates labeling to an injected `labeler` callable, plus
9     a deterministic fake for tests/CI. No concrete model is wired yet ? that's a
10    tracked backlog item (docs/BACKLOG.md), to be chosen alongside the production LLM
11    step so the lineages differ (see the oracle rule below).
12
13    ?? **Oracle rule.** The labeler MUST be a *different model lineage* than whatever
14    production uses. If production becomes the offline ActionFormer segmenter and the
15    oracle is also ActionFormer, they share blind spots and the regression test is
16    blind ? green while both are wrong. `lineage` records the oracle's model family so
17    callers can guard against it matching production (`warn_if_same_lineage`).
18    """
19    import logging
20    from typing import Callable, Dict, List, Optional
21
22    from backend.annotations.base import (
23        AnnotationProvider,
24        annotation_registry,
25        Interval,
26        DONE,
27        FAILED,
28    )
29
30    logger = logging.getLogger(__name__)
31
32    # A labeler turns (video_ref, guideline, sport) into rally intervals.
33    Labeler = Callable[[str, Optional[str], str], List[Interval]]
34

```

```

35
36     def _not_configured(video_ref: str, guideline: Optional[str], sport: str) ->
List[Interval]:
37         """Default labeler: refuses to run until a real model is injected.
38
39         Keeps `AutomatedProvider()` constructible (so the registry can key it as 'auto')
40         while making accidental use of an unconfigured oracle fail loudly rather than
41         silently returning no rallies.
42         """
43         raise NotImplementedError(
44             "AutomatedProvider has no labeler configured. Inject one via "
45             "annotation_registry.get('auto', labeler=<fn>, lineage='<model-family>')."
46             "Choosing the concrete oracle model is a tracked item ? see docs/BACKLOG.md
(GS-4)."
47         )
48
49
50     def warn_if_same_lineage(oracle_lineage: Optional[str], production_lineage:
Optional[str]) -> bool:
51         """Warn (and return True) if the oracle shares production's model lineage.
52
53         Encodes the oracle rule in code, not just docs: a same-lineage oracle can't be
54         trusted to catch production's failures. Callers wiring CI should pass both.
55         """
56         if oracle_lineage and production_lineage and oracle_lineage == production_lineage:
57             logger.warning(
58                 "Oracle lineage '%s' matches production's ? the regression test may be "
59                 "blind to shared failure modes. Use a different model lineage.",
60                 oracle_lineage,
61             )
62             return True
63         return False
64
65
66     class AutomatedProvider(AnnotationProvider):
67         def __init__(self, labeler: Optional[Labeler] = None, lineage: str = "unspecified"):
68             """
69             Args:
70                 labeler: callable producing rally intervals for a video. Defaults to a
71                         stub that raises until configured.
72                 lineage: the oracle's model family (e.g. 'gemini', 'qwen-vl') ? recorded
73                         so callers can enforce the oracle rule via warn_if_same_lineage.
74             """
75             self._labeler = labeler or _not_configured
76             self.lineage = lineage
77             # In-process synchronous job store: job_id -> labeled intervals.
78             self._jobs: Dict[str, List[Interval]] = {}
79
80         @property
81         def name(self) -> str:
82             return "auto"
83
84         def submit(self, video_ref: str, guideline: Optional[str] = None,
85                     sport: str = "badminton") -> str:
86             # Synchronous stub: run the labeler now and stash the result. A real async
87             # oracle (e.g. a queued API job) would kick off here and resolve in poll().
88             intervals = self._labeler(video_ref, guideline, sport)
89             intervals = sorted((float(s), float(e)) for s, e in intervals)
90             job_id = video_ref
91             self._jobs[job_id] = intervals
92             return job_id
93
94         def poll(self, job_id: str) -> str:
95             return DONE if job_id in self._jobs else FAILED
96

```

```

97         def fetch(self, job_id: str) -> List[Interval]:
98             if job_id not in self._jobs:
99                 raise KeyError(f"AutomatedProvider: unknown job '{job_id}' (submit it
first)")
100             return self._jobs[job_id]
101
102
103     annotation_registry.register(AutomatedProvider)
104

```

47. user (2026-06-10T06:05:13.407Z)

```

1     import cv2
2     import time
3     import numpy as np
4     import random
5     from typing import List, Dict, Any, Tuple
6     import logging
7     from backend.infrastructure.database import Database
8     from backend.pipeline.segmenters.base import BasePlaySegmenter, segmenter_registry
9     from backend.results import Result, Success, Failure # standardized contract
10
11     logger = logging.getLogger(__name__)
12
13     class MotionPlaySegmenter(BasePlaySegmenter):
14         @property
15         def name(self) -> str:
16             return "motion"
17
18         @property
19         def description(self) -> str:
20             return "Traditional CV pixel-motion heuristics tracking frame differences."
21
22         def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None, max_frames: int = None) -> List[Dict[str, Any]]:
23             """
24             Processes a video to detect play segments using motion heuristics,
25             populates SQLite with detected segments, and indexes detailed events.
26             """
27             if not player_pool:
28                 player_pool = ["Avi", "Suresh", "Ramesh", "Deepak"]
29
30             cap = cv2.VideoCapture(video_path)
31             if not cap.isOpened():
32                 logger.error(f"Segmenter could not open video: {video_path}")
33                 return Failure(message=f"Could not open video: {video_path}",
is_recoverable=False)
34
35             fps = cap.get(cv2.CAP_PROP_FPS)
36             total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
37             duration = total_frames / fps if fps > 0 else 0
38
39             # Sub-sample settings (5 FPS for analysis)
40             analysis_fps = 5.0
41             frame_interval = max(1, int(fps / analysis_fps))
42
43             idx_cfg = self.config.get("indexing", {})
44             motion_threshold = idx_cfg.get("motion_threshold", 0.08)
45             min_segment_duration = idx_cfg.get("min_segment_duration", 3.0)
46             dead_time_merge_gap = idx_cfg.get("dead_time_merge_gap", 2.0)
47
48             prev_gray = None
49             motion_signals = []
50             timestamps = []
51

```



```

52     frame_idx = 0
53     processed_count = 0
54     t_start = time.time()
55     # Emit a progress line roughly every 5% of the video so long runs are
56     # observable instead of silent (a full 1080p match is ~tens of minutes).
57     progress_step = max(1, total_frames // 20) if total_frames > 0 else 0
58     next_progress = progress_step
59     logger.info(f"[motion] start: {total_frames} frames @ {fps:.1f}fps "
60               f"({duration/60:.1f} min), analyzing every {frame_interval}th
frame")
61     while True:
62         # Skip decoding for intermediate frames using cap.grab() for high
performance
63         for _ in range(frame_interval - 1):
64             if not cap.grab():
65                 break
66             frame_idx += 1
67
68         ret, frame = cap.read()
69         if not ret:
70             break
71
72         # Process frame
73         gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
74         # Apply Gaussian Blur to smooth noise
75         gray = cv2.GaussianBlur(gray, (21, 21), 0)
76
77         t = frame_idx / fps
78         timestamps.append(t)
79
80         if prev_gray is not None:
81             # Frame absolute difference (highly optimized motion tracking)
82             diff = cv2.absdiff(prev_gray, gray)
83             _, thresh = cv2.threshold(diff, 25, 255, cv2.THRESH_BINARY)
84
85             # Dilate to fill gaps
86             thresh = cv2.dilate(thresh, None, iterations=2)
87
88             # Motion ratio
89             motion_ratio = cv2.countNonZero(thresh) / (thresh.shape[0] *
thresh.shape[1])
90             motion_signals.append(motion_ratio)
91         else:
92             motion_signals.append(0.0)
93
94         prev_gray = gray
95         frame_idx += 1
96         processed_count += 1
97
98         if progress_step and frame_idx >= next_progress:
99             pct = 100.0 * frame_idx / total_frames
100             elapsed = time.time() - t_start
101             eta = elapsed / (frame_idx / total_frames) - elapsed if frame_idx else 0
102             logger.info(f"[motion] {pct:5.1f}% ? {frame_idx}/{total_frames} frames,
"
103                       f"{processed_count} analyzed, {elapsed:.0f}s elapsed,
~{eta:.0f}s left")
104             next_progress += progress_step
105
106         if max_frames is not None and processed_count >= max_frames:
107             logger.debug(f"Reached max_frames limit of {max_frames}. Halting indexer
frame parsing.")
108             break
109
110         if frame_idx >= total_frames:

```

```

111         break
112
113     cap.release()
114
115     # Smooth signal using a moving average
116     window_size = 5
117     smoothed = []
118     for i in range(len(motion_signals)):
119         start = max(0, i - window_size // 2)
120         end = min(len(motion_signals), i + window_size // 2 + 1)
121         smoothed.append(sum(motion_signals[start:end]) / (end - start))
122
123     # Detect active play boundaries
124     in_segment = False
125     segment_start = None
126     raw_segments: List[Tuple[float, float]] = []
127
128     for i, t in enumerate(timestamps):
129         signal = smoothed[i]
130         if signal > motion_threshold:
131             if not in_segment:
132                 in_segment = True
133                 segment_start = t
134             else:
135                 if in_segment:
136                     in_segment = False
137                     raw_segments.append((segment_start, t))
138
139     # Handle video ending while in active segment
140     if in_segment:
141         raw_segments.append((segment_start, timestamps[-1]))
142
143     # Post-Processing: 1. Merge segments separated by less than dead_time_merge_gap
144     merged_segments: List[Tuple[float, float]] = []
145     if raw_segments:
146         curr_start, curr_end = raw_segments[0]
147         for next_start, next_end in raw_segments[1:]:
148             if next_start - curr_end < dead_time_merge_gap:
149                 # Merge
150                 curr_end = next_end
151             else:
152                 merged_segments.append((curr_start, curr_end))
153                 curr_start, curr_end = next_start, next_end
154         merged_segments.append((curr_start, curr_end))
155
156     # Post-Processing: 2. Filter segments shorter than min_segment_duration
157     final_segments: List[Tuple[float, float]] = [
158         r for r in merged_segments if (r[1] - r[0]) >= min_segment_duration
159     ]
160
161     # Populate SQLite DB with segments and metadata events
162     segments_data = []
163     for start, end in final_segments:
164         # Assign Server and Receiver randomly from pool
165         server, receiver = random.sample(player_pool, 2)
166
167         # Formulate dynamic rich tags
168         dur = end - start
169         metadata = {
170             "shot_count": random.randint(4, 25),
171             "intensity": "high" if dur > 12 else "medium",
172             "avg_speed_kmh": round(random.uniform(40, 85), 1)
173         }
174
175         segment_id = self.db.add_play_segment(video_id, start, end, server,

```

```

receiver, metadata)
176
177         if segment_id:
178             # Add sub-segment events
179             # A: Serve Event (always at start)
180             serve_foul = random.random() < 0.15 # 15% chance of serve foul
181             self.db.add_event(
182                 segment_id=segment_id,
183                 timestamp=start + 0.5,
184                 event_type="serve",
185                 player=server,
186                 details={"status": "foul" if serve_foul else "valid", "type":
"low_serve" if random.random() > 0.5 else "high_serve"}
187             )
188
189             # B: Smash Events (1 to 3 random smashes per segment)
190             smashes_count = random.randint(0, 3)
191             for _ in range(smashes_count):
192                 smash_time = random.uniform(start + 1.5, end - 1.0)
193                 smasher = random.choice([server, receiver])
194                 self.db.add_event(
195                     segment_id=segment_id,
196                     timestamp=round(smash_time, 2),
197                     event_type="smash",
198                     player=smasher,
199                     details={"speed_kmh": round(random.uniform(120, 250), 1)}
200                 )
201
202             # C: Segment Termination Event (foul/winner)
203             foul_types = ["out_of_bounds", "net_touch", "double_hit", "fault"]
204             is_foul = random.random() > 0.4
205             error_player = random.choice([server, receiver])
206
207             self.db.add_event(
208                 segment_id=segment_id,
209                 timestamp=end - 0.2,
210                 event_type="foul" if is_foul else "winner",
211                 player=error_player,
212                 details={"type": random.choice(foul_types)} if is_foul else
{"shot_type": "drop" if random.random() > 0.5 else "smash"}
213             )
214
215             # Store segment representation to return to caller
216             segments_data.append({
217                 "id": segment_id,
218                 "start_time": start,
219                 "end_time": end,
220                 "duration": dur,
221                 "server": server,
222                 "receiver": receiver,
223                 "metadata": metadata
224             })
225
226         return Success(value=segments_data)
227
228     # Register the segmenter
229     segmenter_registry.register(MotionPlaySegmenter)
230

```

48. user (2026-06-10T06:05:14.667Z)

```

tests\test_base.py:17:         return Failure(message="Simulated failure for test",
is_recoverable=True)
tests\test_base.py:18:         return Success(value=[{"start_time": 0.0, "end_time": 1.0}])
backend\results.py:27:class Success(Generic[T]):

```

```

backend\pipeline\segmenters\motion.py:33:         return Failure(message=f"Could not open
video: {video_path}", is_recoverable=False)
backend\pipeline\segmenters\motion.py:226:         return Success(value=segments_data)

```

49. user (2026-06-10T06:05:25.544Z)

```

1     import subprocess
2     import os
3     import shutil
4     import logging
5
6     logger = logging.getLogger(__name__)
7
8     def get_video_duration(video_path: str) -> float:
9         """Returns the exact duration of a video file in seconds using ffprobe.
10        Returns 0.0 on failure."""
11        try:
12            duration_output = subprocess.check_output([
13                "ffprobe", "-v", "error", "-show_entries", "format=duration",
14                "-of", "default=noprint_wrappers=1:nokey=1", video_path
15            ]).decode().strip()
16            return float(duration_output)
17        except Exception as e:
18            logger.error(f"Could not detect video duration for {video_path}: {e}")
19            return 0.0
20
21    def extract_video_chunk(video_path: str, start_time: float, duration: float,
22                           output_path: str, reencode: bool = False,
23                           scale_width: int = None) -> bool:
24        """Extracts a chunk from a video file.
25
26        Args:
27            reencode: If True, re-encodes with libx264 (slower but frame-accurate).
28                     If False, uses stream-copy (fast but cuts at nearest keyframe).
29                     Use True for short clips where keyframe drift matters.
30            scale_width: If set (requires reencode), downscale to this width (height auto,
31                        even). Useful for high-res sources headed to an API that analyses at
32                        low res anyway ? shrinks uploads with no analysis loss.
33
34        """
35        cmd = [
36            "ffmpeg", "-y",
37            "-ss", str(start_time),
38            "-t", str(duration),
39            "-i", video_path,
40        ]
41        if reencode:
42            if scale_width:
43                cmd += ["-vf", f"scale={int(scale_width)}:-2"]
44            cmd += ["-c:v", "libx264", "-preset", "ultrafast", "-crf", "23", "-c:a", "aac"]
45        else:
46            cmd += ["-c", "copy"]
47        cmd.append(output_path)
48        try:
49            subprocess.run(cmd, check=True, stdout=subprocess.DEVNULL,
50                           stderr=subprocess.DEVNULL)
51            return True
52        except Exception as e:
53            logger.error(f"Failed to extract chunk {start_time}s - {start_time+duration}s:
54            {e}")
55            return False
56

```

50. user (2026-06-10T06:05:25.989Z)

```

1     import math

```

```

2     from typing import Tuple
3
4     def get_center(bbox: Tuple[float, float, float, float]) -> Tuple[float, float]:
5         """Calculate the center (x, y) of a bounding box.
6         bbox format: (x1, y1, x2, y2)
7         """
8         x1, y1, x2, y2 = bbox
9         return ((x1 + x2) / 2.0, (y1 + y2) / 2.0)
10
11    def get_distance(p1: Tuple[float, float], p2: Tuple[float, float]) -> float:
12        """Calculate the Euclidean distance between two points (x, y)."""
13        return math.sqrt((p2[0] - p1[0])**2 + (p2[1] - p1[1])**2)
14
15    def get_velocity(bbox1: Tuple[float, float, float, float], bbox2: Tuple[float, float,
float, float], fps: float) -> float:
16        """Calculate velocity (pixels per second) of a bounding box moving between two
consecutive frames.
17        """
18        c1 = get_center(bbox1)
19        c2 = get_center(bbox2)
20        dist = get_distance(c1, c2)
21        return dist * fps
22
23    def get_velocity_normalised(bbox1: Tuple[float, float, float, float], bbox2:
Tuple[float, float, float, float],
24                                fps: float, frame_width: float) -> float:
25        """Calculate velocity as a fraction of frame width per second.
26
27        This makes the threshold resolution-independent:
28        - 0.15 means the center moved 15% of the frame width in one second.
29        - Works identically on 720p, 1080p, and 4K footage.
30
31        Args:
32            bbox1: Previous bounding box (x1, y1, x2, y2).
33            bbox2: Current bounding box (x1, y1, x2, y2).
34            fps: Effective frames per second (after frame-skipping).
35            frame_width: Width of the video frame in pixels.
36        """
37        raw_velocity = get_velocity(bbox1, bbox2, fps)
38        if frame_width <= 0:
39            return raw_velocity # Fallback to raw pixels if width unknown
40        return raw_velocity / frame_width
41
42    def calculate_iou(boxA: Tuple[float, float, float, float], boxB: Tuple[float, float,
float, float]) -> float:
43        """Calculate the Intersection over Union (IoU) of two bounding boxes.
44        bbox format: (x1, y1, x2, y2)
45        """
46        xA = max(boxA[0], boxB[0])
47        yA = max(boxA[1], boxB[1])
48        xB = min(boxA[2], boxB[2])
49        yB = min(boxA[3], boxB[3])
50
51        interArea = max(0, xB - xA + 1) * max(0, yB - yA + 1)
52        if interArea == 0:
53            return 0.0
54
55        boxAArea = (boxA[2] - boxA[0] + 1) * (boxA[3] - boxA[1] + 1)
56        boxBArea = (boxB[2] - boxB[0] + 1) * (boxB[3] - boxB[1] + 1)
57
58        iou = interArea / float(boxAArea + boxBArea - interArea)
59        return iou
60

```

```

1      """Canonical file-identity hashing for the indexer.
2
3      `video_id` is the MD5 of the *whole* video file. Historically each endpoint
4      (`backend/main.py`, `run_eval.py`, `run_phase3_eval.py`, `run_process_and_stitch.py`)
5      defined its own `get_file_md5` with a different read-chunk size (64KB / 8KB / 1MB).
6      The hash values agreed (same bytes hashed), but the duplicated code was a footgun:
7      any future change to one copy would silently diverge the DB key.
8
9      This module is the single source of truth. Import `compute_video_id` everywhere a
10     video needs to be keyed.
11     """
12
13     from __future__ import annotations
14
15     import hashlib
16
17     # 64 KB matches the value used by the live pipeline (backend/main.py) before this
18     # was centralized. The chunk size does not affect the resulting digest ? it only
19     # bounds memory use while streaming large files.
20     _CHUNK_SIZE = 65536
21
22
23     def compute_video_id(filepath: str) -> str:
24         """Return the MD5 hex digest of the whole file ? the canonical video_id.
25
26         Streams the file in fixed-size chunks so multi-gigabyte recordings never load
27         into memory at once.
28         """
29         hasher = hashlib.md5()
30         with open(filepath, "rb") as f:
31             for chunk in iter(lambda: f.read(_CHUNK_SIZE), b""):
32                 hasher.update(chunk)
33         return hasher.hexdigest()
34

```

52. user (2026-06-10T06:05:26.831Z)

```

1      # Observability report spec for the badminton-highlight-indexer.
2      # The report ENGINE (config language, renderers) lives in the shared `obsreport`
3      # library; this file is the indexer-specific layout + footgun rules. Every
4      # actionable rule cites a README/FAQ section so a reader can self-debug.
5      name: indexer
6      title: Badminton Highlight Indexer ? Video Report
7
8      # Customer here = whoever receives the highlight reel.
9      customer_verdict:
10         pass: "? Your highlights are ready."
11         pass_with_warnings: "? Your highlights are ready. (A couple of minor notes above.)"
12         fail: "?? We couldn't finish generating highlights for this video. Please contact
support."
13
14     sections:
15         - { id: run,          title: "Run",          kind: identity,  audience: both }
16         - { id: video,       title: "Video",         kind: video_meta, audience: both }
17         - { id: highlights,  title: "Highlights",     kind: highlights, audience: customer
18     }
19         - { id: cust_notes,  title: "Things to know", kind: flags,     audience: customer
20     }
21         - { id: stages,     title: "Processing stages", kind: stages,    audience:
technical }
22         - { id: findings,   title: "Findings & warnings", kind: flags,     audience:
technical }
23         - { id: verdict,    title: "Verdict",        kind: verdict,   audience: both }
24
25     rules:

```

```

24 # The classic silent failure: 0 segments because the AI provider no-op'd on a
25 # missing API key ? NOT because the video has no rallies.
26 - code: silent_provider_noop
27 severity: error
28 when:
29   - { path: "stages.ai_handoff.details.api_key_present", op: eq, value: false }
30   - { path: "stages.segmentation.metrics.segment_count", op: eq, value: 0 }
31 message: >-
32   0 segments AND no AI provider API key was set. The segmenter (see Processing
33   stages > segmentation > segmenter_actual) hands its candidate windows to the AI
34   provider, which silently did nothing without a key ? so this is a MISSING
35   CREDENTIAL, not 'the video has no rallies'. Fix: put GEMINI_API_KEY=<key> in a
36   .env file in the project root and re-run (see README#Q7).
37 faq_ref: "README#Q7"
38 customer_message: >-
39   We couldn't generate highlights for this video due to a configuration issue
40   on our side. Please contact support.
41
42 # API key WAS present but an AI segmenter still produced nothing ? likely a
43 # provider error / quota / WSL healthcheck issue rather than a truly empty video.
44 - code: ai_empty_vs_no_rallies
45 severity: warn
46 when:
47   - { path: "stages.ai_handoff.details.ai_based", op: eq, value: true }
48   - { path: "stages.ai_handoff.details.api_key_present", op: eq, value: true }
49   - { path: "stages.segmentation.metrics.segment_count", op: eq, value: 0 }
50 message: >-
51   An AI-based segmenter returned 0 segments even though the API key is set.
52   Empty can mean a provider error, exhausted quota, or (for WASB/TrackNet) a
53   failed WSL healthcheck ? these all return [] silently. Re-run with -u for
54   unbuffered logs and check the provider/WSL output before concluding 'no rallies'.
55 faq_ref: "README#Q7"
56
57 # fps was assumed rather than probed -> velocity thresholds + timing may be wrong.
58 - code: fps_fallback_used
59 severity: warn
60 when:
61   - { path: "video_meta.fps_source", op: eq, value: fallback }
62 message: >-
63   fps was NOT probed from the file ? a fallback default (shown as the fps value)
64   was used. Because rally timing and motion thresholds scale with fps, boundaries
65   and motion sensitivity may be miscalibrated. This usually means validation didn't
66   run (it probes the file) or ffprobe couldn't read r_frame_rate. Confirm the real
67   fps with ffprobe and re-run.
68 faq_ref: "README#Q14"
69
70 # Audio stream missing or disabled. Per VIDEO_RECORDING_GUIDELINES (and ADR-007)
71 # the shuttle "thwack" is a first-class rally cue for recall; recording with audio
72 # on (mic unobstructed, no aggressive noise suppression) is now a hard requirement
73 # for good results with the shuttle-tracking substrates.
74 - code: audio_disabled_or_missing
75 severity: warn
76 when:
77   - { path: "video_meta.audio_present", op: eq, value: false }
78 message: >-
79   No audio stream was detected in the input (or it could not be read). The shuttle
80   "thwack" impact sound is a strong cue that improves rally recall (especially for
81   the WASB/TrackNet detector paths). Record with audio enabled, keep the camera
82   mic unobstructed, and avoid heavy wind/noise-reduction filters that flatten the
83   transient. See VIDEO_RECORDING_GUIDELINES.md (the "Audio (REQUIRED)" row) and
84   AUDIO_RALLY_DETECTION_PLAN.md.
85 faq_ref: "VIDEO_RECORDING_GUIDELINES"
86 customer_message: >-
87   For best highlight results, record with audio on and the mic unobstructed.
88   The shuttle impact sound helps the detector find rallies. Re-record or

```

```

89         re-process after enabling audio (see VIDEO_RECORDING_GUIDELINES.md).
90
91     # The motion baseline fabricates its per-segment metadata/events ? they are not
92     # measured. Warn so nobody trusts them as ground truth.
93     - code: synthetic_motion_metadata
94       severity: warn
95     when:
96       - { path: "stages.segmentation.details.segmenter_actual", op: eq, value: motion }
97     message: >-
98       The 'motion' segmenter was used: its per-segment metadata (intensity, speed,
99       events, server/receiver) are PLACEHOLDER values (unmeasured) ? fabricated by the
100      heuristic, NOT derived from the video ? so a PASS with segments here does not mean
101      the metadata is trustworthy. Use yolo_hybrid/gemini for real shot metadata.
102      faq_ref: "README#Q10"
103
104     # The segmenter that ran differs from config's default -> you may be looking at a
105     # different detector than expected (config/setup/CLI defaults are known to diverge).
106     - code: segmenter_divergence
107       severity: info
108     when:
109       - { path: "stages.segmentation.details.matches_config_default", op: eq, value:
false }
110       - { path: "stages.segmentation.details.segmenter_explicitly_requested", op: eq,
value: false }
111     message: >-
112       The segmenter that ran differs from config.json's default_segmenter, and no
113       explicit --segmenter was given ? so a fallback default kicked in. The indexer's
114       defaults are known to diverge across config.json / setup.py / the CLI; confirm
115       you ran the detector you meant.
116       faq_ref: "README#Q7"
117
118     # Re-processing an already-indexed video replaced its prior segments (INSERT OR
119     # REPLACE cascades). Note it so a changed result isn't mistaken for instability.
120     - code: reingest_replaced_prior
121       severity: info
122     when:
123       - { path: "stages.segmentation.details.was_reingest", op: eq, value: true }
124     message: >-
125       This video was already in the database; re-processing REPLACED its previous
126       segments (videos are keyed by file MD5; re-ingest cascades a delete). Expected
127       when you intend to re-run; just know the earlier results are gone.
128       faq_ref: "README#Q11"
129
130     # A validation check failed -> segmentation was halted. Point the reader at the
131     # report's Validation section + the per-check FAQ.
132     - code: validation_failed
133       severity: error
134     when:
135       - { path: "stages.validation.status", op: eq, value: error }
136     message: >-
137       One or more validation checks failed, so indexing was halted. See the
138       Processing stages > validation block above: each failed check states the
139       reason (relevance -> README#Q2, fps/keyframes -> README#Q5).
140       faq_ref: "README#Q13"
141     customer_message: >-
142       This video couldn't be processed because it didn't pass our quality checks
143       (for example: too blurry, too low resolution, or not a badminton court).
144

```

53. assistant (2026-06-10T06:05:38.278Z)

Now let me verify the key dynamic claims ? nested-config extra-key dropping, import chain weight ? and run the test suite.

54. user (2026-06-10T06:05:42.232Z)


```
wasb key present: False
yolo_prefetch_workers present: False
debug_duration present: False
default_segmenter: wasb_hybrid
top-level report extra kept: {'enabled': False}
```

55. assistant (2026-06-10T06:05:52.127Z)

Confirmed ? the nested-section key dropping is real. Now check the import chain weight and run the test suite.

56. user (2026-06-10T06:06:05.738Z)

```
---exit: 0
```

57. user (2026-06-10T06:06:26.812Z)

```
cv2 4.13.0
torch 2.6.0+cu124
numpy 2.4.6
segmenters pkg OK: ['motion', 'gemini', 'yolo_hybrid', 'tracknet_hybrid', 'wasb_hybrid']
```

58. user (2026-06-10T06:06:45.552Z)

package import OK without torch

59. user (2026-06-10T06:07:12.449Z)

```
=====
Full test suite: C:\Users\avidu\AppData\Local\Programs\Python\Python313\python.exe -m pytest
-q
obsreport (optional reporting dep): available
=====
..... [ 28%]
..... [ 56%]
..... [ 84%]
..... [100%]
256 passed in 24.30s
=====
RESULT: PASS - all tests green
=====
```

60. user (2026-06-10T06:07:25.645Z)

package import FAILS without torch: torch blocked (simulating GPU-less API host)

61. assistant (2026-06-10T06:08:02.393Z)

Test suite green (256 passed). A few final verification greps.

62. user (2026-06-10T06:08:03.987Z)

```
tests\test_export_wasb_segments.py:4:from backend.eval import export_wasb_segments as ex
tests\test_rally_slicer.py:7:from backend.tools.export import basic_stats, segments_to_json,
segments_to_csv, segments_to_edl
```

63. user (2026-06-10T06:08:05.380Z)

```
backend\pipeline\segmenters\gemini.py:60: # A: Check for debug_duration to trim the video
first
backend\pipeline\segmenters\gemini.py:61: debug_duration = self.config.get("indexing",
{ }).get("debug_duration")
backend\pipeline\segmenters\gemini.py:68: if debug_duration:
```

```
backend\pipeline\segmenters\gemini.py:69:         logger.info(f"DEBUG FEATURE: Trimming video
to first {debug_duration} seconds for fast processing...")
backend\pipeline\segmenters\gemini.py:81:         "-t", str(debug_duration),
backend\pipeline\segmenters\yolo_hybrid.py:309:         handoff_padding =
idx_cfg.get("ai_handoff_padding", idx_cfg.get("gemini_handoff_padding", 2.0))
backend\pipeline\segmenters\yolo_hybrid.py:310:         ai_max_workers =
idx_cfg.get("ai_max_workers", idx_cfg.get("gemini_max_workers", 5))
backend\pipeline\segmenters\yolo_hybrid.py:342:         prefetch_workers =
idx_cfg.get("yolo_prefetch_workers", 2)
```

64. user (2026-06-10T06:08:07.339Z)

```
app.js
index.html
style.css
---
__init__.py
__pycache__
test_annotation_providers.py
test_automated_provider.py
test_base.py
test_calibrate_local.py
test_calibrate_wasb.py
test_calibration.py
test_config.py
test_database.py
test_distill_local.py
test_eval_annotations.py
test_eval_batch.py
test_eval_classifier.py
test_eval_harness.py
test_eval_metrics.py
test_eval_partial_labels.py
test_eval_regression.py
test_eval_training.py
test_export_wasb_segments.py
test_gemini.py
test_gemini_refine.py
test_geometry.py
test_hashing.py
test_hit_detector.py
test_import_smoke.py
test_provider_gemini.py
test_providers.py
test_rally_gate.py
test_rally_seq_proto.py
test_rally_slicer.py
test_reporting.py
test_reporting_harvest.py
test_reporting_spec_faq.py
test_sport_profiles.py
test_tracknet_hybrid.py
test_tracknet_runner.py
test_validators_base.py
test_video.py
test_wasb_cache.py
---
test_gemini_api.cpython-311-pytest-9.0.3.pyc
test_gemini_api.cpython-313-pytest-9.0.3.pyc
```

65. user (2026-06-10T06:08:52.932Z)

```
config.json:38:         "yolo_parallel_chunks": false,
config.json:39:         "yolo_chunk_workers": 2,
README.md:134:        "yolo_parallel_chunks": false,
```

```
README.md:135:         "yolo_chunk_workers": 2,  
backend\config\models.py:59:     yolo_parallel_chunks: bool = False  
backend\config\models.py:60:     yolo_chunk_workers: int = 2
```

66. user (2026-06-10T06:08:54.046Z)

No matches found

67. assistant (2026-06-10T06:10:14.913Z)

All evidence gathered. Assembling the structured audit output.

68. user (2026-06-10T06:12:36.861Z)

Structured output provided successfully

69. assistant (2026-06-10T06:12:46.471Z)

Audit complete. Structured output delivered with 20 findings (6 high, 8 medium, 5 low, 1 info) and 8 evidenced strengths.