

claude-windows / 45ea3a91-654d-4972-9cd3-65f35db54caf

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\45ea3a91-654d-4972-9cd3-65f35db54caf\subagents\workflows\wf_2a64bc88-2a8\agent-a498c995a00e4346d.jsonl`  
- SHA-256: `df5b84c23db2cf2e18236963ff3d8d537658ca9200e74027ef289d28e9404f97`  
- Source modified: `2026-06-10T06:11:20+00:00`  
- Imported at: `2026-07-05T16:48:27+00:00`  
- project: `wf_2a64bc88-2a8`  
- session_id: `45ea3a91-654d-4972-9cd3-65f35db54caf`
```

Transcript

1. user (2026-06-10T06:09:07.938Z)

CONTEXT ? two sibling repos on a Windows machine:

- INDEXER: C:/Users/avidu/Projects/badminton-highlight-indexer ? local AI badminton(->racket-sports) highlight indexer + rally detector. FastAPI + SQLite + ffmpeg + GPU CV via WSL; Gemini = paid cloud AI. docs/ tree is rich. Current branch docs/next-steps-multisport (master = main branch).

- COLLECTOR: C:/Users/avidu/Projects/sports-data-collector ? golden-label acquisition/curation/ingestion CLI (system-of-record candidate for training corpus).

Project state (June 2026): scaffolding complete; owned-model roadmap agreed (docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md) but NO platform/owned-model implementation started; next committed platform slice = async job model. Licensing guardrail: MIT/Apache-2.0/BSD only for code+data+weights; paid LLMs are inference-only, never training-data sources. ENVIRONMENT RULES: this dev environment has NO GPU/torch/cv2 ? do NOT attempt to run the video pipeline or import torch/cv2 modules. Test suites are offline/fully-mocked and SAFE to run. You are READ-ONLY: never edit/create/delete repo files; running tests and git read commands is allowed.

QUALITY BAR: cite file paths (and line numbers where possible) for every finding. Report what IS in the code/docs, not what you assume. Be skeptical and concrete. Your final structured output is consumed programmatically by an orchestrator ? return raw data, not prose for a human.

ROLE: adversarial verifier. An auditor ("docs-strategy") claimed the finding below. Try to REFUTE it against the actual files. Default to isReal=false if the evidence does not hold up, the behavior is documented/deliberate (e.g. listed in docs/CODE_MAP.md gotchas or DEFERRED_HARDENING_PLAN.md as a known deferral ? unless the auditor adds genuinely NEW consequence), or the severity is materially overstated.

FINDING: {"title":"CLAUDE.md and docs/README.md never reference the 'authoritative roadmap' (OWNED_MODEL_IMPLEMENTATION_PLAN.md) or NEXT_STEPS.md","severity":"high","area":"doc index / ramp-up path","file":"C:/Users/avidu/Projects/badminton-highlight-indexer/docs/README.md","evidence":"docs/NEXT_STEPS.md:3-4: 'Authoritative roadmap: docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md'. But CLAUDE.md 'Read these first' (lines 7-15) lists only COMMERCIALIZATION/PMF/PLATFORM/DEFERRED_HARDENING/ROADMAP/BACKLOG, and docs/README.md (the doc index) lists neither NEXT_STEPS.md nor OWNED_MODEL_IMPLEMENTATION_PLAN.md anywhere. Repo-wide grep for 'OWNED_MODEL_IMPLEMENTATION_PLAN|NEXT_STEPS' matches only NEXT_STEPS.md itself, the archived study's banner, and quality-iteration archives. CLAUDE.md's 'Current state (June 2026)' (lines 17-21) cites merges only through '#52/#53/#55/#57/#58' though #59-#63 are merged on master (git log: 4d04a73 = #61 owned-model plan, a5477c2 = #62 NEXT_STEPS), and makes no mention of the owned-model plan or Phase-0.", "why_it_matters":"The two entry-point documents (CLAUDE.md, docs/README.md) route a cold-start reader to the 2026-06-07 strategy docs that the 2026-06-09 plan partially supersedes (licensing wording, Gemma-4, Qwen2.5-vs-3, conversational-QA framing) ? and never to the doc that supersedes them. The reader would confidently act on stale strategy."}

Steps: open the cited file(s) and any counter-evidence files; verify quotes/line claims; judge severity honestly. Cheap greps allowed; do NOT run the pipeline. Return isReal, confidence, and a 1-3 sentence verdict (with corrected detail if partially right).

2. user (2026-06-10T06:09:21.887Z)

```
1      # Next steps (live) ? owned-model + multisport
2
3      **Updated:** 2026-06-09 (end of the quality-iteration + multisport-recon session).
**Authoritative roadmap:**
4      `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`. **Latest blow-by-blow:** memory `current-
state.md` / `session-handoff.md`.
5      This file = the prioritized "what to do next," refreshed each session.
6
7      ## The recommendation (one paragraph)
8      Build the **quality-first extraction harness** before anything else: a learned **TAL
head** (ActionFormer/ASFormer, MIT)
9      over trajectory features is the robust path ? it **decouples from the multi-court
"background games" confound** that the
10     unsupervised heuristic (and audio, and spatial court-gating) can't escape (all three
tested + failed; it's a *temporal*
11     rally-vs-non-rally problem). **Framework = ms-swift, base = Qwen3-VL** (Gemma-4
AMBER/bench) come **later** for the
12     conversational/grounding upgrade ? start with the head. **On-device is deferred**
(export-clean ops now so it stays
13     reachable). Everything stays **MIT/Apache-2.0/BSD** (code + data + weights). Bottleneck
= **data + per-sport trajectory
14     detectors**, not model architecture.
15
16     ## Latest findings (2026-06-09)
17     - **6-match badminton LOVO:** learned prototype **0.583** vs heuristic **0.480**
(learned leads since n=3, AUC 0.83?0.92,
18     decoupled from multi-court). The "contrast" metric (in-rally÷dead-time visibility)
predicts heuristic F1 monotonically.
19     - **?? Multisport (Roora pickleball + TT):** labels **ingest cleanly** (`curate convert-
cvat` ? 6 rallies/sport, sane).
20     **WASB-transfer SURPRISE (directional, n=6 each ? TINY):** the badminton shuttle
detector segments **TABLE TENNIS at
21     F1 0.909** out-of-the-box (!) but **pickleball only 0.308**. ? **The per-sport-
detector blocker is SPORT-DEPENDENT:** TT may
22     be servable with WASB short-term; pickleball needs a proper (MIT/Apache) ball
detector. Confirm/deny with more videos.
23     - **Owner field notes (2026-06-09):** camera height/distance **varied on purpose**
(great generalization data); **multi-court
24     adjacent-court SOUND dominates the signal** ? empirically reconfirms audio-foreground
isolation won't work + the
25     robustness-first learned-model path. Multi-court is the realistic (academy)
environment, so this *is* the target distribution.
26
27     ## Prioritized next steps
28     1. **[owner greenlight] Phase-0 (M0.x), $0/local:**
29     - **M0.4 ? Gen-0 TAL harness** (first-class): productionize `rally_seq_proto.py` ?
one-command train?eval?ship-gate** over
30     ActionFormer/ASFormer, reusing `metrics.match_intervals`. **Define the ship-gate
target** (explicit F1@tIoU; floor = beat
31     heuristic LOVO 0.480) ? "very high P/R" must be operationalized, not just "beat the
floor."
32     - **M0.1 ? corpus spine + the per-video event-index contract** (the conversational
seam: extract ? event store ? text-to-SQL).
33     - **M0.3 ? compliance:** purge Gemini-derived training labels
(`training.label_candidates`); WASB **C3** own-weights/clean-swap (gates ship).
34     2. **Multisport thread:**
35     - Ingest the Roora pickleball/TT rally CSVs into the corpus (collector `import-
local`).
36     - **Get more pickleball + TT videos** ? re-run the WASB-transfer eval to confirm/deny
TT?0.91, pickleball?0.31 (n=6 is too few).
37     - Once each sport has ?3 matches, run a **merged prototype LOVO across badminton + TT
(+ pickleball)** ? the "how would a
38     merged model do" experiment. Identify a clean MIT/Apache ball detector for sports
```

```

where WASB doesn't transfer (pickleball).
39     3. **Housekeeping:** **rename repo off "badminton" (multisport); confirm collector =
corpus/event-store system-of-record.
40
41     ## Pointers / artifacts
42     - Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Study:
`docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md`
43     (2026-06-09 banner) · Quality: `docs/archives/quality-iterations/2026-06-multivideo-
lovo/`.
44     - Eval corpus: `output/human_lovo_manifest.json` (badminton n=6). Roora CSVs:
`%TEMP%\roora\{pickleball,tabletennis}_rallies.csv`;
45     WASB-transfer trajectories: `%TEMP%\{pickleball,tabletennis}_traj.csv`. Code:
`backend/eval/{eval_partial_labels,rally_seq_proto,calibrate_local}.py`.
46     - ? No owned-model implementation until owner greenlights a milestone. Committed next
PLATFORM slice = async job model (single-tenant).
47

```

3. user (2026-06-10T06:09:22.169Z)

```

1     # docs/
2
3     Live documentation for the badminton-highlight-indexer.
4
5     ## For new contributors / ramp-up (read these first)
6     - [CODE_MAP.md](CODE_MAP.md) ? Complete navigation map of the codebase and data
contracts.
7     - [COMMERCIALIZATION.md](COMMERCIALIZATION.md) ? Living tracker for the commercial
dimension (licensing, codec/output economics, hosting, vendoring, packaging) + PMF risks and the
quality/data-phase priorities. Supersedes the archived `COMMERCIAL_READINESS_REVIEW.md`.
8     - [PMF_MARKET_RESEARCH.md](PMF_MARKET_RESEARCH.md) ? Cited market study: sizing,
customer segments, competitor pricing, and the recommended beachhead (feeds COMMERCIALIZATION
$6).
9     - [PLATFORM_ARCHITECTURE.md](PLATFORM_ARCHITECTURE.md) ? Strategy for the multi-user
"assistant coach" platform: decoupled layers + pluggable storage/compute (BYO), multi-tenant
data model, orchestration, and a phased roadmap.
10    - [DEFERRED_HARDENING_PLAN.md](DEFERRED_HARDENING_PLAN.md) ? Design (pending approval)
for the larger deferred items: packaging, train/eval guardrails, and the async job model (#16,
the first infra slice toward the platform).
11    - [BACKLOG.md](BACKLOG.md) ? Deferred items with rationale and revisit triggers.
12    - [ROADMAP.md](ROADMAP.md) ? High-level four-phase product direction.
13    - [VIDEO_RECORDING_GUIDELINES.md](VIDEO_RECORDING_GUIDELINES.md) ? Capture best
practices (referenced by reports).
14    - [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md) ? Mental model of the
two-layer detection system.
15    - [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
16
17    ## Practical / ongoing
18    - [TRACKNET_WSL_SETUP.md](TRACKNET_WSL_SETUP.md)
19    - Golden set docs (`GOLDEN_SET_`)
20    - Evaluation harness and calibration drivers (in `backend/eval/`)
21
22    ## Historical documents
23    Completed plans, old research, ADRs, and checkpoints live in [`archives/`](archives/).
24
25    See `archives/README.md` for the umbrella policy and `archives/past_projects/README.md`
for classification of small tech debt / starter projects.
26

```

4. user (2026-06-10T06:09:22.578Z)

```

1     # CLAUDE.md ? project memory / agent handoff
2
3     Local, AI-powered **badminton (? racket-sports) highlight indexer + rally detector**:
4     FastAPI + SQLite + ffmpeg + GPU CV/AI, with paid **Gemini** as the cloud AI. Sibling
5     repos: `sports-data-collector`, a VLC-based `rally-annotator`.

```

```

6
7     ## Read these first
8     - **`docs/README.md`** ? doc index. **`docs/CODE_MAP.md`** ? code navigation + data
9       contracts (the cold-start map; read before touching code).
10    - **Strategy (post-scaffolding)**: **`docs/COMMERCIALIZATION.md`**,
11      `docs/PMF_MARKET_RESEARCH.md`,
12      `docs/PLATFORM_ARCHITECTURE.md`, `docs/DEFERRED_HARDENING_PLAN.md`. Also
13      `docs/ROADMAP.md`
14      (offline-segmenter narrative) and `docs/BACKLOG.md`.
15    - **History**: **`docs/archives/`** ? ADRs (`decisions/DECISIONS.md`), research, and
16      **`checkpoints/`** (latest: `2026-06-strategy-foundation/`). Archive policy: only move
17      **terminal-state** (DONE/REJECTED) work; live docs stay at `docs/` top level.
18
19    ## Current state (June 2026)
20    - **Scaffolding complete**; the **post-scaffolding strategy foundation** is laid and
21      merged
22      (#52/#53/#55/#57/#58). **No platform/owned-model implementation has started.**
23    - **Next slice = async job model** (`DEFERRED_HARDENING_PLAN.md` #16), single-tenant.**
24      **Do not start implementation without explicit owner approval.**
25
26    ## Architecture in one breath
27    - **Pluggable registries** (ABC + `Registry` singleton): segmenters / providers /
28      validators
29      / sports / annotations ? see `backend/pipeline/segmenters/base.py`. *Future:*
30      `StorageBackend` + `ComputeBackend` registries (PLATFORM §3).
31    - `video_id` = MD5 of the file ? `backend/utils/hashing.py::compute_video_id` (one
32      helper).
33    - Typed config = `backend/config/` (Pydantic). DB schema evolves via the `PRAGMA
34      user_version` migration ladder in `backend/infrastructure/database.py`.
35
36    ## Committed guardrails (non-negotiable)
37    - **Paid LLMs (Gemini/OpenAI/Anthropic) = inference / serving / fallback ONLY ? never a
38      training data source** (terms/copyright). Training teachers/bases = **open Apache-2.0
39      only**.
40    - **Owned-model base**: Apache-2.0, video-capable (Qwen-VL / Gemma 4); **reject Gemma
41      3**
42      (viral license). **Exclude minors** from the training pool; per-user training data
43      needs a
44      separate explicit opt-in.
45    - **v1 = rally + highlight + history + correction loop**; **gait/biometrics strictly
46      future**
47      (GDPR Art. 9).
48    - **Generic data schema**; `coach ? athlete` is an **optional overlay**, not baked in.
49
50    ## Workflow conventions
51    - **Branch off `master`**; ONE PR per work item; NO PR stacking.** Don't open a PR unless
52      asked.
53    - Tests: `python run_all_tests.py` (offline, fully mocked). Python 3.8+; `ffmpeg`
54      required.
55    - Never push to the default branch directly; descriptive commit messages.
56    - ?? This dev container lacks GPU / torch / cv2 / numpy ? the **full pipeline can't run
57      here**.
58    - Verify pure-logic paths and hand machine-side checks (real video, GPU/WSL suite) to
59      the owner.
60
61

```

5. user (2026-06-10T06:09:34.205Z)

```

docs\NEXT_STEPS.md:4: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`. **Latest blow-by-blow:** memory
`current-state.md` / `session-handoff.md`.
docs\NEXT_STEPS.md:42:- Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Study:
`docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md`
docs\archives\research\OWNED_MODEL_TRAINING_STUDY.md:12:## Update (2026-06-09, post-
PR-#61-review) ? authoritative roadmap is now `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`
docs\archives\quality-iterations\2026-06-multivideo-

```

lovo\README.md:116: `docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md` (Gen-0 first), gated on n?5 + owner sign-off.
docs\archives\quality-iterations\2026-06-multivideo-lovo\README.md:253: +
`docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md`, after scaffolding).

6. user (2026-06-10T06:09:34.892Z)

13d4573 fix(stitcher): restore f-string quotes dropped in logging conversion (SyntaxError on master) (#63)
a5477c2 docs: next-steps + multisport findings (WASB transfers to table tennis, not pickleball) (#62)
4d04a73 Quality iteration: human-GT multi-video LOVO + learned prototype + multi-court finding + owned-model study/plan (#61)
90fb8b3 Human-GT partial-label eval harness + tuning + iteration doc (#60)
3fd7232 docs: handoff-consistency sweep + strategy checkpoint + CLAUDE.md (#59)
7609593 Merge pull request #58 from avidullu/claude/platform-architecture-research
757842c docs(platform): address review ? v1-vs-vision scope, Hatchet framing, BYO ops burden
bc7b99e docs(platform): generic schema, committed legal guardrail, per-user adapters, rich-metadata stage
380239e docs(platform): add nanochat-style agent-runnable training harness (\$5A)
81838f5 docs(platform): survey training-backend alternatives to Tinker (\$5A)
f83999b docs(platform): drop redundant --- rules that rendered as paragraph underlines
8fe780e docs(commercialization): group dashboard Pending -> In-flight -> Completed + add Priority (P0/P1/P2) column
7e534a3 docs(platform): firm owned-model verdicts + citations (research follow-up)
67410c8 docs(platform): owned generative video-QA model layer + data flywheel
b5833d1 docs: multi-user "assistant coach" platform architecture strategy
master

7. user (2026-06-10T06:10:04.676Z)

docs\EVALUATION.md
docs\TRACKNET_WSL_SETUP.md
docs\GOLDEN_SET_IMPLEMENTATION.md
docs\GOLDEN_SET_PHASE1_VIDEOS.md
docs\HOW_RALLY_DETECTION_WORKS.md
docs\DEFERRED_HARDENING_PLAN.md
docs\PMF_MARKET_RESEARCH.md
docs\ROADMAP.md
docs\COMMERCIALIZATION.md
docs\PLATFORM_ARCHITECTURE.md
docs\BACKLOG.md
docs\GOLDEN_SET_VENDORING.md
docs\README.md
docs\VIDEO_RECORDING_GUIDELINES.md
docs\CODE_MAP.md
docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md
docs\NEXT_STEPS.md

8. user (2026-06-10T06:10:06.264Z)

c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:47: an unscoped new workstream (coordinate labels or detector swap); (5) OWNED_MODEL_IMPLEMENTATION_PLAN + NEXT_STEPS are
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:48: linked from NO entry point (CLAUDE.md "read first" / docs/README index) and NEXT_STEPS is branch-only; (6) README
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:103: authoritative doc =
** docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md ** (full verified decisions in the "(PR review)" checkpoint
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:131: `docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md`, the authoritative roadmap; pushed to PR #61 + a consolidated PR comment reply):**

```
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:203: (`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` DRAFT, 4 decisions pending).
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:225:- **OWNED-MODEL: PLAN-ONLY per owner** ?
`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` written (DRAFT, uncommitted, for
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\session-handoff.md:21:**The authoritative roadmap =
`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`** (rewritten + owner-agreed in #61). Verified
c:\Users\avidu\.claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\session-handoff.md:66:- **ROADMAP: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`**
(authoritative). Study: `docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md`
```

9. user (2026-06-10T06:10:21.063Z)

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
4      only on request). What just landed, what's next, open decisions."
5      metadata:
6          node_type: memory
7          type: project
8          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
9      ---
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-10 (compiler SyntaxError FIX ? **PR #63 MERGED**, squash `master
13     13d4573`; owner approved merge-on-green) .
14     ? **master UNBROKEN again ? fixed the ? SyntaxError from 7fbbeb0** (= the audit's fix
15     chip `task_2b7a19d4` work):
16     restored the f-string quotes on the 7 mangled lines of
17     `backend/pipeline/stitcher/compiler.py`
18     (100/104/156/160/181/191/213; message texts verified against the pre-conversion file at
19     `7fbbeb0`; line 191 =
20     plain string, no placeholders) + fixed line 168 `f"\nSUMMARY]" ? `f"\n[SUMMARY]`. **NEW
21     regression guard
22     `tests/test_import_smoke.py`**: (a) py_compile sweep over every file in backend/
23     (imports nothing ? runs on
24     cv2-free machines), (b) subprocess `import backend.main` with missing cv2/torch/numpy
25     stubbed. Negative check:
26     the broken compiler.py fails the sweep (`SyntaxError: invalid decimal literal`, line
27     100).
28     ? **CI NOW EXISTS ? the audit's no-CI finding is CLOSED** (`github/workflows/ci.yml`,
29     runs on every PR + master
30     push): job 1 **import-smoke** = smoke tests on a GPU-stack-free runner (minimal dep set
31     verified empirically in a
32     fresh venv: pytest/fastapi/uvicorn/pydantic/python-dotenv/**google-genai** ? needed at
33     import time by
34     backend.providers); job 2 **full-suite** = requirements.txt with **CPU-only torch
35     wheels**
36     (`--index-url download.pytorch.org/whl/cpu`) + libgl for opencv ? `python
37     run_all_tests.py` (obsreport absent ?
38     reporting tests skip by design). **Both jobs GREEN on the first run** (27s / 1m29s).
39     Verified post-merge:
40     `import backend.main` OK on master, suite 256 green (254+2).
41     ?? Audit's **coverage-inversion** finding (wasb_hybrid/validators untested) remains open
42     ? CI only guards the
43     import/syntax class + whatever the suite covers.
44
45     **Checkpoint:** 2026-06-10 (DEEP AUDIT session ? ultracode 18-agent audit
46     `wf_2a64bc88-2a8` over indexer+collector:
47     engineering/testing/reliability/docs-consistency/PMF; 3 of 8 auditors
48     completed+verified, 5 rate-limited).
49     - ? **CRITICAL ? master is BROKEN since 7fbbeb0 (2026-06-07 print?logger sweep):**
```

```

33     `backend/pipeline/stitcher/compiler.py` has **7 unquoted f-string lines
(100/104/156/160/181/191/213)** ? SyntaxError ?
34     `import backend.main` fails ? FastAPI server + `/api/compile` dead. Triple-verified +
reproduced first-hand
35     (`py_compile` exit 1). The **254-green suite never noticed** (zero tests import
backend.main / compiler / wasb_hybrid /
36     motion / any concrete validator; **no CI exists at all**). Spawned fix chip
`task_2b7a19d4` (fix + import/compileall
37     smoke test). Coverage is INVERTED vs production risk (tracknet twin 8 tests vs live
wasb_hybrid 0; TrackNetRunner 21
38     vs WasbRunner 1). Suite itself healthy where it exists (254 green / 19.9s / faithful
mocks / low flake risk); "fully
39     mocked runs anywhere" is false on cv2-free machines (bare `import cv2` at collection
time, only obsreport importorskips).
40     - **Doc drift (verified, one-directional ? newest decisions not back-propagated):** (1)
licensing "Apache-2.0 only" still
41     in CLAUDE.md:33 / PLATFORM:269,301 / COMMERCIALIZATION C14 vs ratified
**MIT/Apache/BSD** (a literal CLAUDE.md reader
42     rejects the plan's own MIT ActionFormer/ASFormer heads); (2) **Gemma-4 still
"approved"** in those same docs vs
43     AMBER/bench-only in the plan; (3) **M0.3 purge pointer WRONG** ?
`training.label_candidates` is a pure GT-agnostic fn;
44     the live Gemini-silver training violation = `backend/eval/distill_local.py` (+
calibrate_local thresholds) + gitignored
45     `output/local_rally_model.json`; (4) **C3 escape hatch broken** ? "train-own weights
(Phase-4 roadmap)" cites a phase
46     that contains no such plan and ADR-002 says corpus labels are temporal-only (CAN'T
train a detector) ? own-weights is
47     an unscoped new workstream (coordinate labels or detector swap); (5)
OWNED_MODEL_IMPLEMENTATION_PLAN + NEXT_STEPS are
48     linked from NO entry point (CLAUDE.md "read first" / docs/README index) and NEXT_STEPS
is branch-only; (6) README
49     "yolo_hybrid by default" vs config.json `gemini` vs CODE_MAP "wasb = LIVE default";
(7) ROADMAP still says audio
50     "gated on golden labels" (trigger now satisfied!) vs plan "parked, not in reserve";
(8) Roora multisport CSVs live ONLY
51     in `%TEMP%` (OS-clearable) + n=6 manifest gitignored ? headline
TT-0.909/pickleball-0.308 evidence is one cleanup from
52     loss; (9) TT-on-WASB idea silently walks into the C3 ship-gate ("multiplies per
sport"); (10) ship-gate F1@tIoU still
53     undefined while 0.44/0.54/0.480 floors circulate; repo-rename "ASAP" tracked in no
tracker (BACKLOG lacks it).
54     - **PMF (web-grounded, sources in audit output):** pricing/cost benchmarks VERIFIED
(SwingVision $179.99/yr ~20k paying
55     subs rev $2.75M +12% YoY; PB Vision $99/yr + **$8/hr-of-video partner API**; Gemini
COGS ~$0.30?0.60/hr ? fat margin);
56     but ***"India badminton whitespace" FALSIFIED ? Bangalore is the most contested
badminton-AI city** : Machaxi ($1.5M
57     Rainmatter+Prakash Padukone), VISIST.AI (70+ Bengaluru academies, Padukone School),
Game Theory (Gopichand-backed smart
58     courts w/ auto-highlights), Stupa ($3.3M, BWF-approved, TT?badminton) ? NONE in our
2026-06-07 competitive table.
59     Surviving wedge (narrow, real): **software-only rally-precise indexing of EXISTING
messy multi-court BYO footage**
60     (market solves it only with per-court hardware) sold per-hour as engine/API, possibly
THROUGH those players
61     (PB Vision×PodPlay template). Fastest falsification: C13 interviews (still undone) +
**Gemini-wrapper baseline vs the
62     n=6 LOVO corpus** (if wrapper ? learned 0.583, owned-model is premature) + 1 paid
academy pilot.
63     - **5 auditors PENDING (rate limit, resets ~15:50 IST): eng-pipeline, eng-eval,
reliability, collector, docs-technical.**
64     Resume (cached prefix): `Workflow({scriptPath: "C:\\Users\\avidu\\.claude\\projects\\C
--Users-avidu-Projects-sports-data-
collector\\45ea3a91-654d-4972-9cd3-65f35db54caf\\workflows\\scripts\\deep-project-audit-

```

```
wf_2a64bc88-2a8.js", resumeFromRunId: "wf_2a64bc88-2a8"}})\`.
65     Full audit JSON: `%LOCALAPPDATA%\Temp\claude\C--Users-avidu-Projects-badminton-
highlight-indexer\45ea3a91-654d-4972-9cd3-65f35db54caf\tasks\wf38bovts.output`.
66
67     **Checkpoint:** 2026-06-09 (Roora MULTISPORT recon ? for FRESH next session) . ??
**Pickleball + Tabletennis golden labels
68     arrived; INGESTION = SOLVED, but the "merged model" experiment is BLOCKED on per-sport
trajectory detectors (validates the
69     roadmap blocker).** Owner left a recon-only ask (deliberately did NOT drain context /
spin heavy jobs ? owner wants me fresh).
70     - **Deliverables:** `Annotation Setup/6 June 2026/{pickleball,tabletennis}/*.zip` (3
CVAT export FORMATS each: COCO
```

10. user (2026-06-10T06:11:15.516Z)

Structured output provided successfully

11. assistant (2026-06-10T06:11:20.571Z)

Verification complete.