

claude-windows / 45ea3a91-654d-4972-9cd3-65f35db54caf

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\45ea3a91-654d-4972-9cd3-65f35db54caf\subagents\workflows\wf_2a64bc88-2a8\agent-acac33e88bc67de86.jsonl`
- SHA-256: `7bdc0f5b7a9630fa5ad63ef724b557167bc3af161fa8c9c07a131414eec32203`
- Source modified: `2026-06-10T05:43:16+00:00`
- Imported at: `2026-07-05T16:48:27+00:00`
- project: `wf_2a64bc88-2a8`
- session_id: `45ea3a91-654d-4972-9cd3-65f35db54caf`
```

Transcript

1. user (2026-06-10T05:41:39.828Z)

CONTEXT ? two sibling repos on a Windows machine:

- INDEXER: C:/Users/avidu/Projects/badminton-highlight-indexer ? local AI badminton(->racket-sports) highlight indexer + rally detector. FastAPI + SQLite + ffmpeg + GPU CV via WSL; Gemini = paid cloud AI. docs/ tree is rich. Current branch docs/next-steps-multisport (master = main branch).

- COLLECTOR: C:/Users/avidu/Projects/sports-data-collector ? golden-label acquisition/curation/ingestion CLI (system-of-record candidate for training corpus).

Project state (June 2026): scaffolding complete; owned-model roadmap agreed (docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md) but NO platform/owned-model implementation started; next committed platform slice = async job model. Licensing guardrail: MIT/Apache-2.0/BSD only for code+data+weights; paid LLMs are inference-only, never training-data sources. ENVIRONMENT RULES: this dev environment has NO GPU/torch/cv2 ? do NOT attempt to run the video pipeline or import torch/cv2 modules. Test suites are offline/fully-mocked and SAFE to run. You are READ-ONLY: never edit/create/delete repo files; running tests and git read commands is allowed.

QUALITY BAR: cite file paths (and line numbers where possible) for every finding. Report what IS in the code/docs, not what you assume. Be skeptical and concrete. Your final structured output is consumed programmatically by an orchestrator ? return raw data, not prose for a human.

ROLE: adversarial verifier. An auditor ("docs-strategy") claimed the finding below. Try to REFUTE it against the actual files. Default to isReal=false if the evidence does not hold up, the behavior is documented/deliberate (e.g. listed in docs/CODE_MAP.md gotchas or DEFERRED_HARDENING_PLAN.md as a known deferral ? unless the auditor adds genuinely NEW consequence), or the severity is materially overstated.

FINDING: {"title":"Licensing policy wording split: 'Apache-2.0 only' (CLAUDE.md, PLATFORM, COMMERCIALIZATION, checkpoint) vs ratified 'MIT/Apache-2.0/BSD' (plan, NEXT_STEPS)","severity":"high","area":"licensing guardrails","file":"C:/Users/avidu/Projects/badminton-highlight-indexer/CLAUDE.md:33","evidence":"CLAUDE.md:33 says 'Training teachers/bases = **open Apache-2.0 only**'. PLATFORM_ARCHITECTURE.md:269 says 'Training teachers are open, Apache-2.0 models only (Qwen-VL / InternVL / Gemma 4)' ? internally inconsistent since InternVL is MIT ? and its \$5A heading (line ~301) is 'Apache-2.0 only ? licensing is decisive'; COMMERCIALIZATION.md C14 (line 43) says 'Lock to Apache-2.0'. But OWNED_MODEL_IMPLEMENTATION_PLAN.md:24-25 says 'MIT / Apache-2.0 / BSD only. (Ratifies the prior \"Apache-2.0 only\" to include MIT/BSD.)' with decision table row 3 (line 152) '**owner: yes** ?', and NEXT_STEPS.md:13 says 'Everything stays MIT/Apache-2.0/BSD'. OWNED_MODEL_TRAINING_STUDY.md:131-132 had flagged exactly this ('guardrail literally says Apache-2.0 only but the approved list includes MIT ? ratify'). The ratification landed in the plan but was never back-propagated to CLAUDE.md, PLATFORM, or COMMERCIALIZATION.\"","why_it_matters":"A literal reader of CLAUDE.md (the repo-root agent handoff, marked 'non-negotiable') would reject MIT components ? including ActionFormer/ASFormer (both MIT), which are the plan's PRIMARY Gen-0 TAL heads (OWNED_MODEL_IMPLEMENTATION_PLAN.md:123-124), and InternVL3 (MIT), the designated Qwen fallback. The owner-ratified policy is the opposite of what the most-read doc says."}

Steps: open the cited file(s) and any counter-evidence files; verify quotes/line claims; judge

severity honestly. Cheap greps allowed; do NOT run the pipeline. Return isReal, confidence, and a 1-3 sentence verdict (with corrected detail if partially right).

2. user (2026-06-10T05:41:46.968Z)

```
/ sports / annotations ? see `backend/pipeline/segmenters/base.py`. *Future:*
`StorageBackend` + `ComputeBackend` registries (PLATFORM §3).
- `video_id` = MD5 of the file ? `backend/utls/hashing.py::compute_video_id` (one helper).
- Typed config = `backend/config/` (Pydantic). DB schema evolves via the `PRAGMA
  user_version` migration ladder in `backend/infrastructure/database.py`.
```

Committed guardrails (non-negotiable)

- ****Paid LLMs (Gemini/OpenAI/Anthropic) = inference / serving / fallback ONLY ? never a training data source**** (terms/copyright). Training teachers/bases = ****open Apache-2.0 only****.
- ****Owned-model base:**** Apache-2.0, video-capable (Qwen-VL / Gemma 4); ****reject Gemma 3**** (viral license). ****Exclude minors**** from the training pool; per-user training data needs a separate explicit opt-in.
- ****v1 = rally + highlight + history + correction loop****; ****gait/biometrics strictly future**** (GDPR Art. 9).
- ****Generic data schema****; **`coach ? athlete`** is an ****optional overlay****, not baked in.

3. user (2026-06-10T05:41:48.349Z)

```
> **The moat is *not* the model.** Anyone can download Qwen-VL/Gemma. The defensible
> assets are ** (a) the consented, badminton-specific video?QA dataset** and ** (b) the
> eval harness**. Build and message *those* ? not "we own an LLM."
```

```
> **? Committed guardrail (legal, non-negotiable).** We respect the licensing findings in
> full. **Paid frontier models (Gemini / OpenAI / Anthropic) are used for *inference /
> serving / fallback* only ? never as a *training* data source.** Harvesting their outputs
> to train our owned model would breach their "no competing model" terms (a terms/copyright
> violation), so it is prohibited in both engineering and legal process. **Training teachers
> are open, Apache-2.0 models only** (Qwen-VL / InternVL / Gemma 4). Detail in the
> distillation constraint below.
```

Core principle: AI is a pluggable provider with a paid fallback

The codebase already abstracts the cloud AI behind ``provider_registry` / `AIVideoProvider` (`backend/providers/`). Make that a **first-class platform principle**:`

- Every generative capability routes through the ****provider registry****.
- The ****owned model**** plugs in as a new provider (``owned_vlm``); ****Gemini (paid, via Vertex AI)**** stays registered as the ****fallback**** (failover on error/low-confidence, or as the quality ceiling/teacher reference).
- This lets us ****ship today on the known paid stack**** and ****swap in the owned model when its quality clears a measured bar**** ? with zero architectural churn and the paid path always one config flip away. (Mirrors the COMMERCIALIZATION C4 A/B/C strategy: the owned model is the long-run "(A) eliminate the per-call dependency"; paid is "(C)".)

The model / training layer (new sub-layer, fed by the scaffolding)

...

Data curation	Training	Eval	Serving
(consented user vid + golden set + rally/segment/event annotations ? video-QA pairs)	(fine-tune an open, commercially-usable base VLM via LoRA/QLoRA)	(extend the existing eval harness to QA quality + the regression gate)	(self-host inference as a provider / ComputeBackend; paid fallback)

...

This layer ****reuses every other layer****: ****Storage**** holds the training corpus; ****Data**** supplies labels (the ``annotation`` registry + ``candidate_segments`` + golden set are already the labeling seam); ****Compute/Orchestration**** runs training and inference jobs (even ****BYO-compute**** can serve inference on the tenant's GPU); the ****eval harness**** already in ``backend/eval/`` extends to grade QA quality. ****In other words, the rock-solid scaffolding is exactly what makes the model flywheel possible**** ? which is why scaffolding comes first (see §6).

Base model & training (Apache-2.0 only ? licensing is decisive)

****Recommended near-term base: `Qwen2.5-VL-7B-Instruct`**** ? Apache-2.0, native long-video (dynamic-FPS sampling, temporal mRoPE), mature LoRA tooling. Upgrade path
****Qwen3-VL-4B/8B (Apache-2.0)****; ****Gemma 4 multimodal (Apache-2.0, 2026)**** is a clean Google-stack alternative to evaluate.

- ?? ****Lock to Apache-2.0-clean bases**** so you can fine-tune, ****keep weights private, gate behind subscriptions, and never owe open-sourcing or branding****. (Note Qwen2.5-VL is Apache ****only at 7B/32B**** ? 3B/72B use the custom Qwen license.)
- ? ****Reject for the flagship:**** ****Gemma 3**** ? custom terms with a ****viral derivative clause**** (anything fine-tuned on it, *or even trained on its synthetic outputs*, stays

4. user (2026-06-10T05:41:49.272Z)

| C5 | ****Output codec royalties**** | ****P0**** | ? | Highlights are still encoded ****H.264 + AAC**** (``libx264`` in ``backend/pipeline/stitcher/compiler.py``) ? encode-side AVC patent exposure remains. | Switch output to ****AV1 + Opus**** (royalty-free) on an ****Ada-class GPU (L4/L40)**** for HW encode. See §4. Decision: provision Ada-class GPU? |

| C9 | ****Packaging / distribution**** | ****P0**** | ? | No ``pyproject.toml``; the diagnostic ``setup.py`` shadows the build system. | Design in ``DEFERRED_HARDENING_PLAN.md`` (#7) ? ****awaiting your approval**** before code. |

| C12 | ****SaaS foundations (async/auth/tenancy)**** | ****P1**** | ? | Sync ``/api/process`` blocks on long videos; no auth/tenancy. | Multi-user/platform strategy in ``PLATFORM_ARCHITECTURE.md`` (pluggable storage/compute, tenancy, orchestration); `async slice = `DEFERRED_HARDENING_PLAN.md`` (#16). |

| | ****? IN-FLIGHT (in progress / watch)**** | | | |

| C4 | ****Runtime AI strategy (A/B/C + owned model)**** | ****P0**** | ? | Live = ****C Gemini 2.5 Flash, paid tier****. Long-term = ****own a self-hosted generative video-QA model**** (fine-tuned open VLM) as primary, ****paid Gemini always-ready as fallback**** via the provider registry. | Make "provider registry + paid fallback" a core principle; build the owned model ****after**** the scaffolding/flywheel are solid. See ``PLATFORM_ARCHITECTURE.md`` §5A. |

| C13 | ****Beachhead validation**** | ****P0**** | ? | Market research (``PMF_MARKET_RESEARCH.md``) recommends ****India badminton academies**** as the #1 beachhead. Unvalidated. | Cheapest tests: 10?15 academy/coach discovery interviews; rally-detection P/R benchmark on 100+ real amateur clips; fake-door pricing (B2C ~\$99/yr, B2B academy ~\$600/yr). See §6. |

| C8 | ****Golden-set vendoring spend**** | ****P1**** | ? | ADR-006: ****Roora (Bengaluru)**** selected for an initial paid pilot; licensed/owned video only (no user uploads to vendors). | Confirm quality-per-dollar from the pilot; then GS-3 ``HumanVendorProvider``. See ``GOLDEN_SET_VENDORING.md``. |

| C3 | ****Shuttle-tracking weights provenance**** | ****P1**** | ? ?? | WASB-SBDT ****code is MIT (clear)****; distributed ****checkpoints are academic-data-trained**** ? not covered by the MIT code grant. MonoTrack/YOLO-NAS banned (noncommercial). | Train-own weights (Phase-4 roadmap) ****or**** confirm per-checkpoint terms ?? before commercial deploy. |

| C14 | ****Owned-model base licensing + training-data policy**** | ****P1**** | ? ?? | Owned video-QA model (PLATFORM §5A) must build on ****Apache-2.0-clean, video-capable**** bases ? ****Qwen2.5-VL-7B/32B**** (Apache; 3B/72B are *not*), ****Qwen3-VL-4B/8B****, or ****Gemma 4**** ? to keep weights private, gate, and monetize. ****Reject Gemma 3**** (viral derivative clause ? permanent Google PUP), Llama Vision (naming + 700M-MAU), ``gpt-oss`` (text-only). ****Hard rule: never train on paid-LLM (Gemini/OpenAI/Anthropic) outputs**** (competing-model clause). The moat is the ****consented dataset + eval harness****, not the model. | Lock to Apache-2.0; ****open-model teachers only****; separate explicit ****training opt-in****; ****exclude minors**** from the training pool; per-clip consent ledger. See PLATFORM §5A/§7. |

| C6 | ****Inbound HEVC decode**** | ****P2**** | ? ?? | Decoding users' GoPro H.264/****HEVC**** uploads is unavoidable; HEVC pools are fragmented + highest-uncertainty. | Low/uncertain residual; confirm decode-side exposure with counsel before launch. |

| C7 | ****Hosting unit-economics**** | ****P2**** | ? | Per-video cost is *cheap* for Gemini (~\$0.05 input for a 10-min clip); risk is volume × latency, not unit price. Self-host/offline costs unmeasured. | ****Benchmark \$/video & latency on the target VM**** for A/B/C before pricing. See §5. |

| C10 | ****Value layer (product surface)**** | ****P2**** | ? | Minimal exports shipped (EDL/JSON/CSV + ``basic_stats``, PR #51). | Deepen: per-rally tags/notes that feed golden, per-player/match stats, clip gallery. See §6. |

| C11 | ****Input-quality enforcement**** | ****P2**** | ? | Per-video observability reports surface a customer verdict + footgun flags incl. audio-present (PR #50). | Turn more ``VIDEO_RECORDING_GUIDELINES`` dimensions into validators/report rules. |

| C2 | ****`supervision` ByteTrack pin**** | ****P2**** | ? | Pinned ``<0.30.0`` (deprecation ahead). |

Migrate before lifting the pin (BACKLOG). Not a legal risk, a maintenance one. |
| | **? COMPLETED** | | | |
| C1 | **AGPL / copyleft removal** | ? | ? | ADR-005 shipped: AGPL `ultralytics` YOLOv8 replaced
with **torchvision Faster R-CNN (BSD) + `supervision` ByteTrack (MIT)**. No AGPL code/weights in
the stack. | Done. Keep new deps permissive. |

5. user (2026-06-10T05:41:57.225Z)

contextual questions ("longest rally", "make a clip of all service faults", "rallies over 20
shots").

- **On-device is LATER** (a model-compression goal so power users save upload cost/hassle) ?
definitely on the
roadmap, but **not** the immediate target. Don't over-index the strategy on it now.
- **The moat = the consented dataset + the eval harness, not the weights.**

Hard guardrails (non-negotiable ? "never corrupted")

- **Commercial-clean licensing for EVERY dependency ? code, data, AND model weights: MIT /
Apache-2.0 / BSD only.**
(Ratifies the prior "Apache-2.0 only" to include MIT/BSD.) Reject viral/NC/custom-restrictive:
Gemma 3 (viral),
Gemma 3n (Gemma Terms follow derivatives ? even cleaner reject), Llama 4 (700M-MAU cap),
Commons-Clause repos.
- **Never train on paid-LLM (Gemini/OpenAI/Anthropic) outputs.** Also forbidden: public
grounding-CoT datasets that
were Gemini-generated (e.g. TVG-R1's 56K) ? we must originate our own CoT supervision (a real,
binding cost).
- **Exclude minors; per-user training data needs a separate explicit opt-in; split data by
match.**
- ? **Base-model pretraining provenance is unauditale for *every* open VLM (incl. Qwen3-VL).**
Our "no paid-LLM

- **((ii) WASB C3 ? first-class action item (owner):** the academic-data checkpoint provenance is
a **commercial-ship
blocker** that a clean head does NOT launder. Resolve via **own-trained weights** or a clean
detector swap; **multiplies
per sport.** Carry as a tracked blocker that **gates SHIP** (not Gen-0 research).

M0.4 ? Gen-0 TAL harness (**FIRST-CLASS ACTION ITEM** per owner) · *greenlight item*

Productionize `rally\_seq\_proto.py` ? a **one-command train?eval?ship-gate** harness over
**ActionFormer (MIT, 1:1
rally=instance)** and/or **ASFormer (MIT, TAS)** (use `sj-li/MS-TCN2` if MS-TCN++, *not* the
Commons-Clause repo),
reusing `metrics.match\_intervals`. **Define the ship-gate target up front** (open item): a FLOOR
of "beat the honest n=6
heuristic LOVO **0.480**" is necessary but not sufficient ? set an explicit **F1@tIoU** target
for "very high P/R" (e.g.
F1@tIoU=0.5 for highlights, tighter for precise rally cuts) before declaring victory. Trains on
the GPU/WSL box.

Decisions / status

#	item	status
1	Greenlight Phase-0 milestone(s) ? rec: **M0.1 corpus+event-index**, then **M0.4 harness** awaiting owner	
2	WASB C3 ? first-class action item (own-weights / clean swap) **owner: yes** ?	
3	Ratify licensing **MIT / Apache-2.0 / BSD** (code+data+weights) **owner: yes** ?	
4	Collector = corpus + event-store system-of-record **owner: confirm** ?	
5	Rename repo off "badminton" **owner action** ?	
6	Conversational-QA: build seam now, punt model **rec accepted** ?	

Open risks / unknowns (none block Phase-0 research)

Budget/cost to reach Gen-2 ****unquantified**** (largest risk) · ship-gate "very high P/R" target undefined · per-sport clean

6. user (2026-06-10T05:41:58.674Z)

Exit code 2

Next steps (live) ? owned-model + multisport

****Updated:**** 2026-06-09 (end of the quality-iteration + multisport-recon session).
****Authoritative roadmap:****
`docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md`. ****Latest blow-by-blow:**** memory `current-state.md` / `session-handoff.md`.
This file = the prioritized "what to do next," refreshed each session.

The recommendation (one paragraph)

Build the ****quality-first extraction harness**** before anything else: a learned ****TAL head**** (ActionFormer/ASFormer, MIT) over trajectory features is the robust path ? it ****decouples from the multi-court "background games" confound**** that the unsupervised heuristic (and audio, and spatial court-gating) can't escape (all three tested + failed; it's a ***temporal*** rally-vs-non-rally problem). ****Framework = ms-swift, base = Qwen3-VL**** (Gemma-4 AMBER/bench) come ****later**** for the conversational/grounding upgrade ? start with the head. ****On-device is deferred**** (export-clean ops now so it stays reachable). Everything stays ****MIT/Apache-2.0/BSD**** (code + data + weights). Bottleneck = ****data + per-sport trajectory detectors****, not model architecture.

Latest findings (2026-06-09)

- ****6-match badminton LOVO:**** learned prototype ****0.583**** vs heuristic ****0.480**** (learned leads since n=3, AUC 0.83?0.92, decoupled from multi-court). The "contrast" metric (in-rally+dead-time visibility) predicts heuristic F1 monotonically.
- ****?? Multisport (Rooru pickleball + TT):**** labels ****ingest cleanly**** (`curate convert-cvat` ? 6 rallies/sport, sane).
****WASB-transfer SURPRISE (directional, n=6 each ? TINY):**** the badminton shuttle detector segments ****TABLE TENNIS at F1 0.909**** out-of-the-box (!) but ****pickleball only 0.308****. ? ****The per-sport-detector blocker is SPORT-DEPENDENT:**** TT may be servable with WASB short-term; pickleball needs a proper (MIT/Apache) ball detector. Confirm/deny with more videos.
- ****Owner field notes (2026-06-09):**** camera height/distance ****varied on purpose**** (great generalization data); ****multi-court adjacent-court SOUND dominates the signal**** ? empirically reconfirms audio-foreground isolation won't work + the robustness-first learned-model path. Multi-court is the realistic (academy) environment, so this ***is*** the target distribution.

sed: can't read C:/Users/avidu/Projects/badminton-highlight-indexer/docs/OWNED_MODEL_TRAINING_STUDY.md: No such file or directory

7. user (2026-06-10T05:42:00.088Z)

C:\Users\avidu\Projects\badminton-highlight-indexer\CLAUDE.md:33: training data source** (terms/copyright). Training teachers/bases = ****open Apache-2.0 only****.
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PLATFORM_ARCHITECTURE.md:301:### Base model & training (Apache-2.0 only ? licensing is decisive)
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:24:- ****Commercial-clean licensing for EVERY dependency ? code, data, AND model weights: MIT / Apache-2.0 / BSD only.****
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:25:

(Ratifies the prior "Apache-2.0 only" to include MIT/BSD.) Reject viral/NC/custom-restrictive: Gemma 3 (viral),
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:44:
InternVL3 (MIT/Apache) = clean fallback if a Qwen blocker appears. ? base swap is cheap
*only within
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:152:| 3 | Ratify licensing **MIT / Apache-2.0 / BSD** (code+data+weights) | **owner: yes** ? |
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\NEXT_STEPS.md:13:reachable). Everything stays **MIT/Apache-2.0/BSD** (code + data + weights). Bottleneck = **data + per-sport trajectory
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\NEXT_STEPS.md:22: be servable with WASB short-term; pickleball needs a proper (MIT/Apache) ball detector. Confirm/deny with more videos.
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\NEXT_STEPS.md:38: merged model do" experiment. Identify a clean MIT/Apache ball detector for sports where WASB doesn't transfer (pickleball).
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\archives\checkpoints\2026-06-strategy-foundation\README.md:24:- **Paid frontier models (Gemini/OpenAI/Anthropic): inference / serving / fallback only ? never a training data source** (competing-model terms). Teachers/bases = open Apache-2.0 only.
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\archives\research\OWNED_MODEL_TRAINING_STUDY.md:98: Gemma 4 E2B/E4B/12B-Unified (Apache-2.0 + native video/audio), InternVL3 ?38B (MIT/Apache), ASFormer/
C:\Users\avidu\Projects\badminton-highlight-indexer\docs\archives\research\OWNED_MODEL_TRAINING_STUDY.md:131:- **Policy ratifications needed:** guardrail literally says "Apache-2.0 only" but the approved list includes MIT ?

8. user (2026-06-10T05:42:23.240Z)

No matches found

9. user (2026-06-10T05:42:23.904Z)

c:\Users\avidu\claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md:69: PRIVATE PROPRIETARY COMMERCIAL software. **MIT / Apache-2.0 / BSD only.** Ratify (was "Apache-2.0 only"). Enforce religiously.
c:\Users\avidu\claude\projects\c--Users-avidu-Projects-badminton-highlight-indexer\memory\session-handoff.md:37:- **Licensing RATIFIED: MIT / Apache-2.0 / BSD only** for code + data + WEIGHTS ("never corrupted"). ? **Gemini-silver

10. user (2026-06-10T05:42:31.025Z)

No matches found

11. user (2026-06-10T05:42:32.804Z)

- **Base (when we reach the VLM): Qwen3-VL (Apache-2.0, uniform across sizes)** ? long-video context, purpose-built
timestamp emission (second-level localization), multi-turn video chat. Re-confirm the exact size's LICENSE
pre-release. **Gemma-4 stays AMBER and bench-only** ? its Apache grant is likely but its Prohibited-Use posture is
unconfirmed (counsel needed), *and* its video/timestamp capability is disputed; not the primary extractor.
InternVL3 (MIT/Apache) = clean fallback if a Qwen blocker appears. ? base swap is cheap
*only within
timestamp-capable Apache bases* (timestamp data-format + grounding-CoT are base-specific).
- **START WITH THE TAL HEAD, NOT THE VLM** (the decisive call): fine-tuned VLMs **miss strict temporal boundaries**
(best RL-post-trained video VLM ? R@0.7 50% on open-domain benchmarks; "coarse regions, not accurate boundaries").
Boundary accuracy IS the product. The head (`rally\_seq\_proto.py`) is already de-risked on our corpus (held-out
===

revision is verified, given the Gemma-3 rejection. Its timestamp-emission recipe is less documented than Qwen's.

- **Existing distilled model was trained on Gemini silver labels** ? the (correct) eval-only policy now **forbids** this; that pipeline must be purged of Gemini-derived training labels to comply.
- **Policy ratifications needed:** guardrail literally says "Apache-2.0 only" but the approved list includes MIT ? ratify "Apache-2.0/MIT". The MAU-cap rejections (72B/78B) may be over-conservative at our scale.
- **On-device privacy claim is unbenchmarked** ? no tokens/sec for a quantized 4B VLM on 8 GB Turing at 1080p/60fps. **Measure before relying on it.**
- **Amplification is load-bearing for Gen-2** ? 128 raw rallies can't reach the ~5k?20k clip?QA band alone; if augmentation + QA-multiplexing + pseudo-labels underdeliver, Option B is under-fed. Amateur 60fps may need more

12. user (2026-06-10T05:43:09.210Z)

Structured output provided successfully

13. assistant (2026-06-10T05:43:16.748Z)

Verification complete. The finding is real: the "Apache-2.0 only" wording in CLAUDE.md:33, PLATFORM_ARCHITECTURE.md:269/301, COMMERCIALIZATION.md:43, and the 2026-06 checkpoint README contradicts the owner-ratified "MIT/Apache-2.0/BSD" policy in OWNED_MODEL_IMPLEMENTATION_PLAN.md:24-25/152 and NEXT_STEPS.md:13, with no tracked task to reconcile them.