

claude-windows / da34752f-c664-4262-9602-f5d4cfcc6db4

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\da34752f-c664-4262-9602-f5d4cfcc6db4\subagents\workflows\wf_c5f5bca5-d2a\agent-a761b94af2b3baacd.jsonl`
- SHA-256: `a4bb7db2f299ff3c3e73445e60b36c81ef236f025039e06b66d6e73988fd7b4b`
- Source modified: `2026-06-12T10:45:10+00:00`
- Imported at: `2026-07-05T16:48:22+00:00`
- project: `wf_c5f5bca5-d2a`
- session_id: `da34752f-c664-4262-9602-f5d4cfcc6db4`
```

Transcript

1. user (2026-06-12T10:43:40.346Z)

Repo: C:\Users\avidu\Projects\sports-data-collector (Windows). Review commit fd530b5 on branch feat/prep-annotation-proxy (diff vs master: run `git diff master...feat/prep-annotation-proxy` in that repo, or read the files directly).

Changed files:

- src/core/proxy.py (NEW - probe/encode/verify/hash/manifest engine)
- src/cli/curate.py (new prep_annotation_proxy method + argparse subparser + dispatch)
- tests/test_proxy.py (NEW - 22 offline tests)
- docs/INGESTION_PIPELINE.md (step 0 + 2 FAQ entries)

Purpose (ADR-002 in docs/DECISIONS.md): re-encode a source video into a downscaled "annotation proxy" for an annotation vendor, while preserving the EXACT frame count / fps / timeline so frame N on the proxy == frame N on the original (the vendor's CVAT labels are recomputed as seconds=frame/fps by src/parsers/cvat/rally_cvat.py). Invariants: never change fps, never trim, verify parity after encode (delete proxy on failure), record provenance (source+proxy MD5) in a manifest CSV. Suite is 159 green; a live ffmpeg-8.1.1 E2E smoke passed.

A reviewer claims the following defect. ADVERSARIALLY VERIFY it by reading the actual code/files ? try to REFUTE it. Default to refute if the scenario cannot actually occur or the behavior is acceptable/intended for this tool's scope.

CLAIM [medium] INGESTION_PIPELINE.md onboarding FAQ and top diagram skip the new step 0 (C:\Users\avidu\Projects\sports-data-collector\docs\INGESTION_PIPELINE.md:7-23, 112-116)
The heading was updated to 'The five commands (in order)' and step 0 added, but (a) the 'How do I onboard a brand-new video?' FAQ ? the operational checklist someone will actually follow for a new video ? still reads 'Get fps into the intake manifest' then jumps straight to convert-cvat, never mentioning prep-annotation-proxy or 'share the PROXY, process the ORIGINAL'; this is the exact workflow ADR-002 calls 'close to a precondition for Phase B', so omitting it here makes the doc actively misleading; (b) the ASCII diagram at the top still begins at 'vendor delivery' with no step-0 box/arrow (PROXY -> vendor), so the diagram and the 'five commands' list disagree about where the pipeline starts.

Proposed fix: Add a first bullet to the onboarding FAQ: 'Run prep-annotation-proxy on the original; share the proxy (and its drive link) with the vendor; fps + video_id go into the intake manifest.' Add a small 'original --prep-annotation-proxy--> proxy --> vendor' leg (or a one-line note) to the diagram.

2. user (2026-06-12T10:43:47.590Z)

```
feat/prep-annotation-proxy
fd530b5 feat(curate): prep-annotation-proxy - frame-identical vendor handoff re-encode (ADR-002)
4fb2d6f docs: ADR-002 - vendor annotation proxies generated in-repo (curate prep-annotation-proxy) (#21)
a2d08f8 docs(cvat): update detect_issues padding rationale to config-driven values (#20)
```

3. user (2026-06-12T10:43:51.507Z)

```
1      # Rally Ingestion Pipeline ? End-to-End FAQ
2
3      How a vendor's rally annotations become training/eval ground truth, and how we
4      validate each delivery against our conventions. This is the operational companion
5      to the annotation spec (`docs/annotation/ROORA_BRIEF.md` + `ANNOTATION_GUIDE.md`).
6
7      ```
8      vendor delivery          OUR side (collector)          modeling
(indexer)
9      ?????????????????? convert ?????????????????? validate ??????????????????
10     ? CVAT export ? ?????????????? ? canonical CSV ? ?????????? ? report .md ?
11     ? (boxes + attrs)? ? (per rally, ? ? report .json ?
12     ? OR ? ? ? decimal secs)? ? review .csv ??? human
verifies
13     ? canonical CSV ? ?????????????????? ??????????????????
flagged rallies
14     ?????????????????? ? ?
15     ? import-local (validated) ?
16     ? golden labels
17     ?????????????????? export ??????????????????
18     ? SQLite (rally ? ?????????? ? <id>.csv ?
start,end (indexer GT)
19     ? tables) ? ? manifest.json?
20     ?????????????????? ??????????????????
21     ?
22     slice clips / calibrate /
train
23     ```
24
25     ## The five commands (in order)
26
27     ```bash
28     # 0. BEFORE handoff: re-encode the source into a frame-identical annotation proxy
29     # (ADR-002 ? downsampled + scrub-friendly so the vendor can work at frame step 1;
30     # same frame count/fps/timeline, parity-verified, provenance recorded).
31     # Share the PROXY with the vendor; everything downstream keeps using the ORIGINAL.
32     python -m src.cli.curate prep-annotation-proxy --video-path match.mp4
33
34     # 1. Normalise a CVAT export ? our canonical per-rally CSV (skip if the vendor sent a
CSV)
35     python -m src.cli.curate convert-cvat --input job.xml \
36         --video-path match.mp4 --out canonical.csv --issues issues.txt
37
38     # 2. Validate the delivery ? report + annotatable review sheet (the QA gate)
39     python -m src.cli.curate validate-delivery --input job.xml \
40         --video-path match.mp4 --video-id pilot_badminton_01 --out-dir data/validation
41     # ...or validate a CSV directly:
42     python -m src.cli.curate validate-delivery --csv canonical.csv --video-id
pilot_badminton_01
43
44     # 3. After human review, import the (corrected) CSV into the DB
45     python -m src.cli.curate import-local --sport badminton \
46         --video-path match.mp4 --annotations-path canonical.csv --fps 59.94
47
48     # 4. Export indexer-ready ground truth (start,end CSVs + manifest)
49     python -m src.cli.curate export --out-dir data/eval_export
50     ```
51
52     ---
53
54     ## FAQ
55
56     ### What delivery formats do we accept?
```

```

57 Both. A CVAT export (image- or track-mode XML ? rally fields carried as box
58 attributes) is normalised by `convert-cvat`; an already-canonical CSV is taken
59 as-is. Vendors may keep working in CVAT ? we do not require them to hand-author
60 CSVs.
61
62 ### What is the "canonical" CSV?
63 One row per rally:
64 `rally_number, start_mmss, end_mmss, start_time, end_time, shots_count, ending_reason,
last_shot_outcome, score_before, score_after, notes`.
65 `start_time`/`end_time` are decimal seconds; `import-local` reads those.
66
67 ### How are timestamps handled? (the fps thing)
68 We recompute each rally's interval from its frame range × the true fps
69 (seconds = frame / fps). A vendor's CVAT time attributes are ignored ? a
70 step-sampled job that assumes 30 fps writes systematically wrong times. fps comes
71 from --fps or ffprobe of --video-path. mm:ss is also accepted everywhere.
72
73 ### What does validation check?
74 | code | severity | meaning |
75 |---|---|---|
76 | `overlap` | blocker | two rallies' intervals overlap (end[N] > start[N+1]) |
77 | `zero_duration` | blocker | end <= start (single-sample / stray box) |
78 | `frame_out_of_range` | high | a box frame is past the video's last frame |
79 | `shot_density` | medium | long rally with implausibly few shots ? likely an inflated
span (often hides a merged rally) |
80 | `possible_missed_rally` | medium | a long unlabeled gap between two rallies ? check
for an omitted rally |
81 | `ending_reason_vocab` | medium | ending_reason outside the agreed set |
82 | `outcome_vocab` | low | last_shot_outcome outside {in,net,out,n_a} |
83
84 ### Blocker vs. review ? what's the gate?
85 validate-delivery exits non-zero if any blocker exists (so it can gate a bulk
86 loop / CI). Non-blocker issues don't stop the gate but populate the review sheet
87 for the human pass. import-local independently rejects overlaps/zero-duration, so
88 bad ground truth can't slip into the DB.
89
90 ### How are missed rallies caught? (the worst error)
91 Two complementary checks, because a miss shows up two ways:
92 1. Clean omission (a real hole in the timeline) ? possible_missed_rally (gap
heuristic).
93 2. Merged/inflated (the missed rally got swallowed into a neighbour's box, leaving
94 no gap) ? shot_density flags the over-long, under-shot rally, and the human finds
95 the second rally inside it.
96 Neither is a guarantee ? they surface regions to review, they don't certify
97 completeness. (A future enhancement: a motion second-opinion over the raw video.)
98
99 ### What's the review sheet?
100 <id>_review.csv ? one row per rally with auto_flags (which checks fired) plus
101 blank VERIFIED / true_start_mmss / true_end_mmss / true_ending_reason /
102 true_last_shot_outcome / reviewer_comment columns. Open it next to the video
103 (timestamp overlay on) and fill the flagged rows. This is the human gate that turns
104 a delivery into trusted golden labels.
105
106 ### Why two ending fields (ending_reason + last_shot_outcome)?
107 They're orthogonal: ending_reason is why the point ended (winner / forced_error /
108 unforced_error / service_fault / let / other); last_shot_outcome is where the
109 last shot went (in / net / out / n_a). A shuttle that flew long is
110 forced_error/unforced_error + out ? both kept, no information lost.
111
112 ### How do I onboard a brand-new video?
113 1. Get fps into the intake manifest (we provide it to the vendor).
114 2. convert-cvat (if CVAT) ? validate-delivery ? read the report, fill the review
115 sheet for flagged rallies ? correct the CSV ? import-local ? export.
116 3. The same suite runs for every video, so quality stays consistent at scale.
117

```

```

118     ### How do I pull a single rally out to eyeball it?
119     The slicer (indexer):
120     ```bash
121     python -m backend.tools.rally_slicer --video match.mp4 \
122         --labels golden_labels.csv --rallies 7,16 --out-dir output/rally_clips
123     ```
124     Clips are padded to match the highlight compiler (0.5s lead / 1.0s tail).
125
126     ### Does coarse frame sampling hurt boundaries?
127     It depends on the consumer. For highlight compilation, ~0.5s sampling slack is
128     absorbed by the compiler's clip padding, so validate-delivery doesn't flag sampling
129     density. For golden training/eval labels the bar is  $\pm 0.3s$  (Brief §4), and the
130     pickleball/table-tennis pilot showed a 30-frame sampling grid at 59.94 fps puts every
131     boundary  $\approx 0.5s$  off ? structurally out of spec. That's why videos are handed off as
132     annotation proxies (step 0): light enough to annotate at frame step 1, with the
133     first/last box of each rally on the exact contact/dead frame.
134
135     ### What's an annotation proxy, and why not annotate the original?
136     A 4K/59.94 original is too heavy to frame-step in CVAT ? at Phase-B scale (~30-minute
137     videos, >100k frames each) it's not realistic at all, which is what pushed the pilot
138     vendor onto the coarse sampling grid above. prep-annotation-proxy re-encodes the
139     source downsampled (?1080p) with dense keyframes for snappy scrubbing while preserving
140     the exact frame count, fps, and timeline (verified after the encode; a failed proxy
141     is deleted). Frame N on the proxy IS frame N on the original, so returned labels apply
142     to the original unchanged. Provenance (source + proxy MD5) is appended to
143     data/annotation_proxy_manifest.csv ? commit it. The proxy is ONLY for annotation:
144     import-local, export, and the indexer (which keys videos by file MD5) keep using
145     the original. Full rationale: ADR-002 in docs/DECISIONS.md.
146
147     ### Will a non-commercial license block export?
148     The video is kept in the DB regardless of license, but the default export is
149     commercial-only ? a non-commercial source (e.g. CC-BY-NC-SA) is excluded
150     from the eval/training set unless you pass --include-noncommercial. So if a video
151     "goes missing" downstream, check its is_commercial flag first. (The observability
152     report raises a license_noncommercial warning for exactly this case.)
153
154     ### What is the observability report? (*.report.json / *.report.md / *.summary.md)
155     Every validate-delivery, convert-cvat, and import-local run now also writes a
156     per-video observability report (next to the video for import-local; into the
157     --out-dir / next to the canonical CSV otherwise):
158     - *.*<id>.report.json* ? machine-readable: every stage, metric, and finding. Diff
159       two runs to see what changed between deliveries.
160     - *.*<id>.report.md* ? the technical report: parse / validation / registration
161       stages plus findings (one per delivery issue, each citing this FAQ). Read this
162       to debug a delivery.
163     - *.*<id>.summary.md* ? a delivery summary for the vendor / data team: the
164       rally count and coverage plus the plain-language list of rallies to review ? no
165       internal codes or severities.
166     This is in addition to the existing <id>_report.md / _report.json / _review.csv
167     (the human review sheet), which are unchanged.
168

```

4. user (2026-06-12T10:44:05.082Z)

```

diff --git a/docs/INGESTION_PIPELINE.md b/docs/INGESTION_PIPELINE.md
index 3120872..a4f9cf5 100644
--- a/docs/INGESTION_PIPELINE.md
+++ b/docs/INGESTION_PIPELINE.md
@@ -22,9 +22,15 @@ to the annotation spec (docs/annotation/ROORA_BRIEF.md +
ANNOTATION_GUIDE.md

                                     slice clips / calibrate / train

...

-## The four commands (in order)
+## The five commands (in order)

```

```

```bash
+# 0. BEFORE handoff: re-encode the source into a frame-identical annotation proxy
+# (ADR-002 ? downscaled + scrub-friendly so the vendor can work at frame step 1;
+# same frame count/fps/timeline, parity-verified, provenance recorded).
+# Share the PROXY with the vendor; everything downstream keeps using the ORIGINAL.
+python -m src.cli.curate prep-annotation-proxy --video-path match.mp4
+
1. Normalise a CVAT export ? our canonical per-rally CSV (skip if the vendor sent a CSV)
python -m src.cli.curate convert-cvat --input job.xml \
--video-path match.mp4 --out canonical.csv --issues issues.txt
@@ -118,8 +124,25 @@ python -m backend.tools.rally_slicer --video match.mp4 \
Clips are padded to match the highlight compiler (0.5s lead / 1.0s tail).

Does coarse frame sampling hurt boundaries?
-Not in practice: ~0.5s sampling slack is absorbed by the highlight compiler's clip
-padding, so we don't flag sampling density.
+It depends on the consumer. For **highlight compilation**, ~0.5s sampling slack is
+absorbed by the compiler's clip padding, so `validate-delivery` doesn't flag sampling
+density. For **golden training/eval labels** the bar is $\pm 0.3s$ (Brief §4), and the
+pickleball/table-tennis pilot showed a 30-frame sampling grid at 59.94 fps puts *every*
+boundary ?0.5s off ? structurally out of spec. That's why videos are handed off as
+annotation **proxies** (step 0): light enough to annotate at frame step 1, with the
+first/last box of each rally on the exact contact/dead frame.
+
+### What's an annotation proxy, and why not annotate the original?
+A 4K/59.94 original is too heavy to frame-step in CVAT ? at Phase-B scale (~30-minute
+videos, >100k frames each) it's not realistic at all, which is what pushed the pilot
+vendor onto the coarse sampling grid above. `prep-annotation-proxy` re-encodes the
+source downscaled (?1080p) with dense keyframes for snappy scrubbing while preserving
+the **exact frame count, fps, and timeline** (verified after the encode; a failed proxy
+is deleted). Frame N on the proxy IS frame N on the original, so returned labels apply
+to the original unchanged. Provenance (source + proxy MD5) is appended to
+`data/annotation_proxy_manifest.csv` ? commit it. The proxy is ONLY for annotation:
+`import-local`, `export`, and the indexer (which keys videos by file MD5) keep using
+the original. Full rationale: ADR-002 in `docs/DECISIONS.md`.

Will a non-commercial license block export?
The video is **kept** in the DB regardless of license, but the default `export` is

```

## 5. user (2026-06-12T10:44:05.212Z)

```

88-
89-**Phase B makes this untenable, not just suboptimal:** ~12?15 videos **per sport** at
90-~30 min each (Brief §3). A 30-minute 4K/59.94 file is **>100k frames**; frame-accurate
91-annotation (frame-step 1) on 4K originals at that scale is not realistic for the vendor.
92-A lightweight **annotation proxy** (downscaled, scrub-friendly re-encode of the identical
93:timeline) is therefore close to a **precondition for Phase B**, not an optimization.
94-
95-The pipeline is proxy-friendly by construction: annotations join by the human-assigned
96-`video_id` (not file MD5), and `convert-cvat` recomputes all times as `frame / fps`,
97-ignoring vendor time attributes. The Brief's "do not re-encode" rule binds the *vendor*
98-(timeline + checksum identity of the file they annotate), not us ? provided downstream

```

## 6. user (2026-06-12T10:44:38.208Z)

```

docs/INGESTION_PIPELINE.md | 29 ++++
src/cli/curate.py | 147 +++++
src/core/proxy.py | 242 +++++
tests/test_proxy.py | 309 +++++
4 files changed, 724 insertions(+), 3 deletions(-)

```

## 7. user (2026-06-12T10:45:04.492Z)

Structured output provided successfully

## 8. assistant (2026-06-12T10:45:10.128Z)

Verification complete. The claim is real: the onboarding FAQ (lines 112-116) and the top ASCII diagram in `C:\Users\avidu\Projects\sports-data-collector\docs\INGESTION\_PIPELINE.md` were both left untouched by the branch diff and omit step 0, while the "five commands" list and two new FAQ entries in the same file document it ? an internal inconsistency, with the ADR-002 "precondition for Phase B" quote confirmed at `docs/DECISIONS.md:93`. Rated low (not medium) severity because the correct workflow is prominently stated at the top of the same one-page doc.