

claude-windows / 7d6bbb83-42c0-459b-9e23-e14e96d4cb2d

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\subagents\workflows\wf_1dc2acd6-052\agent-ad7968cddaff50d85.jsonl`  
- SHA-256: `58230ce8648bfeba4ed3e24076743a74a0df0f0bc21214b264f7225590b2b345`  
- Source modified: `2026-06-13T20:24:41+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `wf_1dc2acd6-052`  
- session_id: `7d6bbb83-42c0-459b-9e23-e14e96d4cb2d`
```

Transcript

1. user (2026-06-13T20:19:37.466Z)

You are an ADVERSARIAL fact-checker. Theme: "Temporal action detection/segmentation + boundary refinement / post-processing".

Here are a teammate's survey findings (JSON):

```
{  
  "theme": "Temporal action detection/segmentation + boundary refinement / post-processing",  
  "our_pattern": "We segment a long static-camera badminton match into rally [start,end] intervals via: (a) multi-scope \"producers\" that emit time-aligned signals ? frame-scope detectors (frozen WASB/TrackNet shuttle tracker, person detector, audio thwack detector, proposed per-frame shuttle-state classifier), window-scope analyzers (per-candidate-window aggregated invariant features: net-crossings, service-fault halt, person two-sidedness), and session-scope analyzers (warm-up/practice span detection); (b) a \"signal not verdict\" rule: every cue is a (value, confidence) feature column, consumed by a tiny logistic-regression scorer that decides which candidate windows are real rallies ? no hand-coded if/else in the decision path; (c) detection-decision separation: expensive frame detection persisted once, cheap windowing+scoring is pure offline re-scoring; (d) RECONCILERS ? a report-first post-hoc pass that flags where the DECISION contradicts persisted EVIDENCE (shuttle inactive before segment end, zero net-crossings, over-merged segment with internal dead-time, loose boundary slack) and proposes idempotent, fixed-precedence, A/B-gated corrections; (e) an experiment harness landing cues flag-gated/default-OFF, graded with leave-one-video-out (group-by-match) on scarce vendor-labeled golden intervals using temporal IoU, promoted only on measured lift.",  
  "key_papers": [  
    {  
      "title": "Temporal Action Segmentation: An Analysis of Modern Techniques",  
      "authors": "Ding, Guodong; Sener, Fadime; Yao, Angela",  
      "year": "2022",  
      "url_or_id": "arXiv:2210.10352",  
      "what_it_does": "Survey of temporal action segmentation (TAS): defines the canonical task (per-frame labeling of long untrimmed video into contiguous action segments), the dominant frame-classification + multi-stage-refinement paradigm, the over-segmentation error problem, and the standard metrics (frame accuracy, segmental edit score, F1@{10,25,50}). Surveys full/weak/unsupervised supervision regimes.",  
      "venue": "IEEE TPAMI / arXiv",  
      "overlap": "strong",  
      "overlap_explanation": "Our task IS temporal action segmentation specialized to a 2-class (rally / dead-time) problem on a long timeline. This survey is the authoritative source for the terminology, metrics, and the over-segmentation framing our reconciler exists to fix. We should cite it and adopt its vocabulary instead of inventing our own ('producers', 'reconcilers')."},  
    {  
      "title": "MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation",  
      "authors": "Farha, Yazan Abu; Gall, Juergen",  
      "year": "2019",  
      "url_or_id": "arXiv:1903.01945 (CVPR 2019)",  
      "what_it_does": "Seminal TAS method: stacks multiple refinement stages of dilated temporal convolutions over frame features; each stage refines the previous stage's full prediction. Introduces a smoothing (truncated-MSE) loss specifically to penalize over-segmentation.",  
      "venue": "CVPR 2019",  
      "overlap": "partial",  
      "overlap_explanation": "Same high-level pattern ? features -> initial prediction -> refinement stage(s) that suppress over-segmentation. Differs in mechanism: MS-TCN refinement is a learned in-model neural stage trained on many labels; our reconciler is an external, report-first, rule+evidence corrector tuned for a ~6-match scarce-label regime. The shared concept is 'a refinement stage exists precisely to fix over-segmentation'; the smoothing-loss idea is adoptable in spirit (penalize fragmentation)."},  
    {  
      "title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries (ASRF)",  
      "authors": "Ishikawa, Yuchi; Kasai,
```

Seito; Aoki, Yoshimitsu; Kataoka, Hirokatsu", "year": "2021", "url_or_id": "arXiv:2007.06866 (WACV 2021)", "what_it_does": "Decomposes segmentation into two complementary branches: an Action Segmentation Branch (frame classification) and a Boundary Regression Branch; the regressed boundaries REFINES the classifier output, reducing over-segmentation (up to +16.1% F1, +13.7% edit). Explicitly a boundary-aware refinement framework.", "venue": "WACV 2021", "overlap": "strong", "overlap_explanation": "This is the closest published analog to our reconciler's BOUNDARY work: a dedicated boundary signal used to post-correct a segmentation. Our 'tighten boundary slack to serve-contact' and 'truncate segment whose shuttle went inactive before its end' are exactly boundary-regression-style refinement, except we use explicit physical evidence (serve contact, shuttle inactivity) as the boundary cue rather than a learned BRB. Strong overlap on the WHAT; our framing is the rule/evidence variant of the same idea."}, {"title": "Activity Grammars for Temporal Action Segmentation", "authors": "Gong, Dayoung; Lee, Joonseok; Jung, Deunsol; Kwak, Suha; Cho, Minsu", "year": "2023", "url_or_id": "arXiv:2312.04266 (NeurIPS 2023)", "what_it_does": "Induces a context-free grammar of valid action sequences and applies an efficient Viterbi-style parser as an inference-time POST-PROCESSING layer over any base segmenter, correcting structurally-invalid frame-level predictions into grammar-consistent sequences. Improves over-segmentation and interpretability on Breakfast/50Salads.", "venue": "NeurIPS 2023", "overlap": "strong", "overlap_explanation": "This is the single closest match to the RECONCILER concept: a model-agnostic, inference-time corrector that enforces STRUCTURAL PRIORS on a base segmenter's output. Our reconciler rules ('rally must contain ≥ 1 net-crossing', 'a rally cannot extend past shuttle inactivity', 'point/dead-time cadence vs continuous warm-up') are exactly a hand-specified activity grammar / structural prior. We are reinventing grammar-constrained refinement with bespoke rules; we should cite this and consider framing reconciler rules as an explicit grammar."}, {"title": "Post-Processing Temporal Action Detection (GAP: Gaussian Approximated Post-processing)", "authors": "Nag, Sauradip; Zhu, Xiatian; Song, Yi-Zhe; Xiang, Tao", "year": "2023", "url_or_id": "arXiv:2211.14924 (CVPR 2023)", "what_it_does": "A model-agnostic post-processing method that fixes temporal-quantization boundary error by modeling action start/end with Gaussian distributions and inferring boundaries at sub-snippet resolution via Taylor expansion ? applied to stored detections without retraining.", "venue": "CVPR 2023", "overlap": "partial", "overlap_explanation": "Validates our 'persist once, re-score/refine forever' principle: a model-agnostic post-processing stage applied to frozen detections to sharpen boundaries. Different mechanism (Gaussian sub-snippet boundary inference vs evidence-driven rules) and different error (quantization vs decision-vs-evidence contradiction), but the architectural framing ? detection frozen, refinement is cheap post-hoc ? is the same and worth citing as precedent for the design."}, {"title": "BMN: Boundary-Matching Network for Temporal Action Proposal Generation", "authors": "Lin, Tianwei; Liu, Xiao; Li, Xin; Ding, Errui; Wen, Shilei", "year": "2019", "url_or_id": "arXiv:1907.09702 (ICCV 2019)", "what_it_does": "Generates dense temporal action PROPOSALS (candidate [start,end] windows) and scores each with a confidence via a boundary-matching mechanism; canonical 'proposal-then-classify/score' temporal action localization pipeline. Builds on BSN (Lin et al. 2018, arXiv:1806.02964).", "venue": "ICCV 2019", "overlap": "partial", "overlap_explanation": "Our 'candidate windows + learned scorer' IS the proposal-generate-then-score paradigm of temporal action localization (TAL), of which BSN/BMN are the canonical instances. We differ by generating candidates from a physics signal (shuttle activity) rather than a learned proposal net, and scoring with logistic regression on interpretable features rather than a deep confidence head. Same skeleton, simpler/scarcer-label-appropriate instantiation."}, {"title": "Action Spotting using Dense Detection Anchors Revisited (SoccerNet 2022)", "authors": "Soares, Joao V. B.; Shah, Avijit; Biswas, Topojoy", "year": "2022", "url_or_id": "arXiv:2206.07846", "what_it_does": "Sports-video temporal localization: operates on PRE-EXTRACTED frame features (frozen backbone), predicts dense per-anchor detection confidence + a fine-grained temporal displacement to sharpen localization, optimizing the 'tight' (small-tolerance) average-mAP. Winner of SoccerNet Action Spotting 2022.", "venue": "SoccerNet/CVSports workshop 2022", "overlap": "partial", "overlap_explanation": "Direct sports analog of our detection/decision split: heavy features computed once and frozen, a light head decides events, and a temporal-displacement refinement sharpens boundaries. Their 'tight average-mAP at small temporal tolerance' is the metric the field uses for boundary tightness ? directly relevant to grading our boundary-slack reconciler beyond a single tIoU threshold."}, {"title": "Service and End of Rally Detection in Badminton Videos", "authors": "(Springer ACIVS/IbPRIA chapter, 2024)", "year": "2024", "url_or_id": "DOI:10.1007/978-3-031-69073-0_15", "what_it_does": "Directly the same task as ours: detects rally start (service) and rally end in badminton match video using shuttle/player motion cues (per search-result snippets; full text authenticated).", "venue": "Springer LNCS (2024)", "overlap": "partial", "overlap_explanation": "Same problem (rally start/end segmentation in badminton) and similar cues (shuttle motion, service), so it is the most direct task-level prior art and should be cited as such. Overlap on TASK is strong;

overlap on METHOD is unverified (could not read full text) ? I mark it partial and flag that the multi-scope-producers + reconciler + LOVO-harness architecture appears to be our differentiator vs a single end-to-end detector. Confirm method before claiming novelty over it.},{\"title\":\"A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton Videos\",\"authors\":\"(SN Computer Science, Springer, 2025)\",\"year\":\"2025\",\"url_or_id\":\"DOI:10.1007/s42979-025-03935-0\",\"what_it_does\":\"Detects rally start via BiLSTM/CNN on service, and rally END via a gradient/motion rule: rally ends when $\geq 80\%$ of consecutive frames show no shuttle motion. Reports $\sim 95\%$ accuracy across 187 rallies (per search-result snippets; full text auth-gated).\",\"venue\":\"SN Computer Science (2025)\",\"overlap\":\"partial\",\"overlap_explanation\":\"Their rally-end rule ('shuttle no-motion over consecutive frames') is essentially our 'shuttle went inactive' reconciler/state cue, and their service-detection is our service-fault/serve-contact cue ? strong overlap on the SIGNALS. The key difference is they hard-code these as if/else verdicts; our explicit thesis is to demote them to feature columns for a learned scorer ('signal not verdict'). Cite as evidence the cues are sound, and as a direct counterexample to the special-casing we are deliberately avoiding.\"}],\"established_terms\":[\"Temporal Action Segmentation (TAS) ? the per-frame-label-into-contiguous-segments task; ours is binary (rally vs dead-time)\",\"Temporal Action Localization / Detection (TAL/TAD) ? proposal-then-score of [start,end] intervals; matches our candidate-windows + scorer\",\"Action proposal generation (BSN/BMN) ? our 'candidate windows'\",\"Over-segmentation error ? the canonical name for the failure mode our reconciler's 'split/merge' rules target\",\"Boundary regression / boundary refinement (ASRF) ? our boundary-slack-to-serve-contact correction\",\"Smoothing loss ? the standard in-training anti-over-segmentation device (MS-TCN)\",\"Refinement stage / multi-stage refinement ? the field's term for what our reconciler is, externalized\",\"Model-agnostic post-processing (GAP) ? our 'persist once, re-score/refine forever' framing\",\"Activity grammar / grammar-constrained / constrained decoding ? our reconciler rules formalized as structural priors\",\"Action spotting + 'tight' average-mAP at small temporal tolerances (SoccerNet) ? the sports-video metric for boundary precision\",\"Segmental F1@{10,25,50}, segmental edit score, frame accuracy/MoF (Lea et al. 2017) ? the metrics the field reports alongside tIoU\"],\"techniques_to_adopt\":[{\"technique\":\"Report and track segmental edit score and F1@{10,25,50} (plus mean-over-frames) alongside our temporal-IoU metric.\",\"why\":\"These are the field-standard TAS metrics and, unlike a single tIoU threshold, the edit score and F1@k directly quantify over-segmentation (fragmentation/spurious-boundary errors) ? exactly what the reconciler's split/merge rules target. Reporting them makes our lift claims legible and comparable to the literature, and gives the reconciler a metric that rewards fixing fragmentation, not just overlap.\"}, {\"technique\":\"Temporal Action Segmentation: An Analysis of Modern Techniques (Ding et al. 2022); MS-TCN (Farha & Gall 2019)\",\"effort\":\"low\"}, {\"technique\":\"Add a 'tight' / small-tolerance localization metric (average-mAP across multiple tIoU thresholds incl. tight ones, e.g. SoccerNet-style) to grade boundary tightness specifically.\",\"why\":\"Our serve-contact boundary-slack reconciler is a boundary-precision intervention; a single loose tIoU threshold can hide whether boundaries actually tightened. The tight-average-mAP metric was designed to be sensitive to exactly these small temporal errors.\"}, {\"technique\":\"Action Spotting using Dense Detection Anchors Revisited (Soares et al. 2022)\",\"effort\":\"low\"}, {\"technique\":\"Frame the reconciler's structural rules as an explicit 'activity grammar' / constrained-decoding layer and resolve them globally (Viterbi-style) rather than as independent local edits.\",\"why\":\"Activity Grammars shows that enforcing valid-sequence structure as a model-agnostic inference-time parser both reduces over-segmentation and is interpretable ? the same goal as our reconciler. A global parse avoids the precedence/conflict headaches of stacking independent local rules and gives a principled, citable formalism for 'rally must contain a net-crossing', 'no extension past shuttle inactivity', 'warm-up cadence differs from point cadence'.\"}, {\"technique\":\"Activity Grammars for Temporal Action Segmentation (Gong et al. 2023)\",\"effort\":\"medium\"}, {\"technique\":\"Adopt the explicit boundary-regression decomposition: keep a separate boundary signal (serve contact / shuttle-inactivity onset) and use it to refine the scorer's segment ends, mirroring ASRF's ASB+BRB split.\",\"why\":\"ASRF shows a dedicated boundary branch refining a separate classification branch yields large F1/edit gains and is the established cure for over-segmentation. Our cues already provide physical boundary evidence; structuring the pipeline as 'classify window' + 'regress/refine boundary from evidence' aligns us with a proven design and names what we are doing.\"}, {\"technique\":\"Alleviating Over-segmentation Errors by Detecting Action Boundaries / ASRF (Ishikawa et al. 2021)\",\"effort\":\"medium\"}, {\"technique\":\"Add a soft over-segmentation penalty to the scorer's objective or post-filter (the spirit of MS-TCN's smoothing loss), not just per-window classification accuracy.\",\"why\":\"Our logistic-regression scorer optimizes per-window labels and is blind to fragmentation across adjacent windows; a smoothing/merge prior (penalize transitions / encourage contiguous same-label runs) directly attacks over-segmentation, which is the dominant TAS error mode.\"}, {\"technique\":\"MS-TCN (Farha & Gall 2019)\",\"effort\":\"medium\"}, {\"technique\":\"Treat the whole refinement step as model-agnostic post-

processing over frozen detections and A/B it as such (which we already do) ? cite GAP/Activity-Grammars as precedent and borrow their evaluation discipline (measure lift of the post-processor in isolation, on stored outputs).", "why": "GAP and Activity Grammars are exactly 'a refinement applied to frozen model outputs, measured for standalone lift'. This both validates our 'detection ? decision / persist once' architecture and gives us a citable methodological template for the reconciler's A/B-before-auto-apply gate.", "from_paper": "Post-Processing Temporal Action Detection / GAP (Nag et al. 2023)", "effort": "low"}], "pitfalls": ["Over-segmentation is THE dominant error mode in this field, and frame/window accuracy hides it. A scorer optimized for per-window accuracy can look good while producing fragmented or over-merged rallies. Mitigation: report edit score + F1@k and add an explicit anti-fragmentation prior (MS-TCN smoothing-loss spirit; Ding et al. 2022).", "Stacked independent local correction rules (truncate / drop / split / tighten) with a fixed precedence can conflict, oscillate, or be non-idempotent across orderings ? the very reason the field moved to a single global parse (Viterbi over a grammar). Risk of brittle rule interactions; mitigate by resolving structural constraints jointly (Activity Grammars, Gong et al. 2023).", "Boundary/quantization error: candidate windows at a coarse stride introduce temporal-quantization error that no amount of scoring fixes; boundaries need sub-window refinement (GAP, Nag et al. 2023). A loose tIoU threshold will mask whether boundaries are actually tight.", "LOVO leakage even when grouping by match: with only ~6 matches, the same venue/court/players/broadcast style recur, so 'leave-one-video-out' still leaks domain (court color, camera framing, player pair). The TAS literature reports large domain shift precisely because train/test differ by subject/scene; with 6 matches, per-match LOVO variance will be huge and a logistic-regression scorer can overfit a handful of cues to match-specific quirks. Mitigate: report per-fold variance, not just mean; treat absolute numbers as low-confidence.", "Tuning reconciler rule thresholds (e.g. 'inactive for N frames', '>=80% no-motion') on the same scarce golden set used to grade them is implicit fitting on the test set ? the badminton rally-detection papers' hand-tuned no-motion thresholds are exactly this risk. Keep rule-threshold selection inside the LOVO loop, never on the full golden set.", "Calibration/feature-confidence drift in late fusion: each producer emits an independent (value, confidence); these confidences are not mutually calibrated, so a logistic regression can be dominated by an over-confident but unreliable cue (the audio cue's known low discrimination is a live example). Calibrate per-cue confidences before fusion. (Late-fusion miscalibration is a documented multi-modal TAL issue; see audio-augmented TAL fusion work, arXiv:2106.14118.)", "Report-first is good, but 'A/B-measured lift before auto-apply' on ~6 matches risks promoting a correction whose measured lift is within noise. Require lift to exceed per-fold variance, and prefer corrections that improve edit/F1 (fragmentation) not just mean tIoU."], "verdict": "partial-overlap", "verdict_rationale": "The skeleton of our design is well-established and we should cite rather than reinvent it: 'candidate windows + learned scorer' is textbook proposal-then-score Temporal Action Localization (BSN/BMN); 'refinement stage to fix over-segmentation' is the defining theme of Temporal Action Segmentation (MS-TCN, ASRF); and 'model-agnostic post-processing over frozen detections to correct/sharpen output' is exactly GAP and Activity Grammars. Our reconciler in particular is a close re-derivation of grammar-constrained / boundary-regression refinement (Activity Grammars 2023, ASRF 2021) ? its rules are a hand-built activity grammar, so we should adopt that framing and its global-parse machinery rather than a stack of ad-hoc local edits. What is genuinely less-charted (novel-combination, not a clean gap) is the SPECIFIC assembly for a scarce-label (~6 match), interval-only, static-camera racket-sports setting: physics-derived cues as feature columns, a deliberately tiny logistic-regression scorer instead of a deep net, the report-first/A/B-gated/human-in-the-loop discipline around the corrector, and the multi-scope (frame/window/session) producer taxonomy with a warm-up/practice session-level analyzer. Direct task analogs exist in badminton (Springer 2024 service/end-of-rally; SN-CS 2025 BiLSTM+gradient), but they hard-code cues as verdicts ? our 'signal not verdict, learn the weights' stance is a defensible contrast, not a reinvention. Net: the architecture is partial-overlap (cite TAS/TAL/refinement literature and adopt its metrics, grammar/boundary-refinement framing, and over-segmentation cures); the scarce-label, interpretable, governance-disciplined instantiation is the contribution.", "search_queries_used": ["temporal action segmentation survey MS-TCN over-segmentation error post-processing refinement", "temporal action localization proposal generation boundary refinement BMN BSN survey", "ASFormer DiffAct temporal action segmentation boundary smoothing refinement transformer", "badminton rally segmentation video detection shuttlecock trajectory rally start end", "ASRF Alleviating Over-segmentation Errors Action Segmentation Refinement Framework boundary regression", "tennis broadcast video rally point segmentation play break detection dead time", "temporal action detection post-processing soft-NMS Gaussian boundary refinement late fusion calibration", "multi-modal fusion temporal action localization feature-level late fusion sports video few labels logistic regression", "SoccerNet action spotting temporal localization features lightweight classifier dense anchors challenge", "temporal action segmentation evaluation metrics F1@k edit score segmental over-segmentation Lea"]

2017","constraint-aware decoding temporal action segmentation rule-based correction grammar prior inference time","leave-one-video-out cross-validation temporal action segmentation small dataset leakage group split sports"]}]}

Your job is to be skeptical and protect against two failure modes: (1) FABRICATED citations, (2) OVERCLAIMED novelty.

DO, using WebSearch/WebFetch (load via ToolSearch "select:WebSearch,WebFetch" if needed):

1. For each cited paper, search its EXACT title (+authors) to confirm it is a REAL paper. Mark `verified_real` yes/no/uncertain and flag any that don't resolve in `fabrication_flags`. Be strict: a plausible-sounding title that you cannot find is a fabrication flag.
2. For every "novel-combination" or "genuine-gap" verdict, actively try to REFUTE it: run fresh searches for the CLOSEST prior art that would downgrade the verdict toward "reinventing"/"partial-overlap". Default to assuming prior art exists until your searches come up empty. List any `closer_prior_art` you find (with real ids).
3. Give a `corrected_verdict` + rationale. Add `extra_course_corrections` if your searches surfaced adopt/avoid notes.

Return the structured object.

2. assistant (2026-06-13T20:19:40.512Z)

I'll fact-check this survey adversarially. Let me start by loading the web tools and verifying each cited paper.

3. assistant (2026-06-13T20:19:43.597Z)

Now let me verify each cited paper. I'll search for them in parallel batches, starting with the first set of core TAS/TAL papers.

4. user (2026-06-13T20:19:51.828Z)

Web search results for query: "MS-TCN Multi-Stage Temporal Convolutional Network Action Segmentation Farha Gall CVPR 2019 arXiv 1903.01945"

Links: [{"title":"MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation","url":"https://openaccess.thecvf.com/content_CVPR_2019/papers/Abu_Farha_MS-TCN_Multi-Stage_Temporal_Convolutional_Network_for_Action_Segmentation_CVPR_2019_paper.pdf"}, {"title":"MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation","url":"https://arxiv.org/pdf/1903.01945"}, {"title":"MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation - ADS","url":"https://ui.adsabs.harvard.edu/abs/2019arXiv190301945A/abstract"}, {"title":"MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation | IEEE Conference Publication | IEEE Xplore","url":"https://ieeexplore.ieee.org/document/8953830/"}, {"title":"GitHub - sj-li/MS-TCN2: MS-TCN++: Multi-Stage Temporal Convolutional Network for Action Segmentation (TPAMI 2020) · GitHub","url":"https://github.com/sj-li/MS-TCN2"}, {"title":"Temporal Action Segmentation: An Analysis of Modern Techniques","url":"https://arxiv.org/pdf/2210.10352"}, {"title":"GitHub - yabufarha/ms-tcn · GitHub","url":"https://github.com/yabufarha/ms-tcn"}, {"title":"MS-TCN++: Multi-Stage Temporal Convolutional Network for Action Segmentation","url":"https://arxiv.org/pdf/2006.09220"}, {"title":"[1903.01945] MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation","url":"https://arxiv.org/abs/1903.01945"}, {"title":"[PDF] MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation | Semantic Scholar","url":"https://www.semanticscholar.org/paper/MS-TCN:-Multi-Stage-Temporal-Convolutional-Network-Farha-Gall/b4673e744d0ded47fe6df3b6314f79a41359578b"}]

I found comprehensive information about the MS-TCN paper you're searching for. Here's what I found:

MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation

****Authors & Publication:**** The paper is by Yazan Abu Farha and Juergen Gall from the University of Bonn, Germany, and it was published at IEEE Conference on Computer Vision and Pattern Recognition (CVPR) in 2019. The arXiv identifier you mentioned (1903.01945) is correct.

****Main Contribution:**** The paper introduces a multi-stage architecture for the temporal action segmentation task, where each stage features a set of dilated temporal convolutions to generate an initial prediction that is refined by the next one.

****Training & Loss:**** The architecture is trained using a combination of a classification loss and a proposed smoothing loss that penalizes over-segmentation errors.

****Results:**** The model achieves state-of-the-art results on three challenging datasets: 50Salads, Georgia Tech Egocentric Activities (GTEA), and the Breakfast dataset.

****Key Resources:****

- Official paper PDF: https://openaccess.thecvf.com/content_CVPR_2019/papers/Abu_Farha_MS-TCN_Multi-Stage_Temporal_Convolutional_Network_for_Action_Segmentation_CVPR_2019_paper.pdf
- arXiv version: <https://arxiv.org/pdf/1903.01945>
- Official GitHub repository: <https://github.com/yabufarha/ms-tcn>

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

5. user (2026-06-13T20:19:52.241Z)

Web search results for query: "Temporal Action Segmentation An Analysis of Modern Techniques Ding Sener Yao 2022 arXiv 2210.10352"

Links: [{"title": "[2210.10352] Temporal Action Segmentation: An Analysis of Modern Techniques", "url": "https://arxiv.org/abs/2210.10352"}, {"title": "Cost-Sensitive Learning for Long-Tailed Temporal Action Segmentation", "url": "https://arxiv.org/html/2503.18358"}, {"title": "?????Temporal Action Segmentation: An Analysis of Modern Techniques | PPT", "url": "https://www.slideshare.net/ttamaki/temporal-action-segmentation-an-analysis-of-modern-techniques"}, {"title": "Spatial-temporal graph transformer network for skeleton-based temporal action segmentation | Multimedia Tools and Applications | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s11042-023-17276-8"}, {"title": "End-to-End Streaming Video Temporal Action Segmentation with Reinforce Learning", "url": "https://arxiv.org/pdf/2309.15683"}, {"title": "[PDF] Temporal Action Segmentation: An Analysis of Modern Techniques | Semantic Scholar", "url": "https://www.semanticscholar.org/paper/8b8dd714d1db923cc3eabc33e21bca88aeeb77ac"}, {"title": "Temporal Action Segmentation: An Analysis of Modern Technique", "url": "https://www.researchgate.net/publication/364437795_Temporal_Action_Segmentation_An_Analysis_of_Modern_Technique"}, {"title": "Temporal Action Segmentation: An Analysis of Modern Techniques - ADS", "url": "https://ui.adsabs.harvard.edu/abs/2022arXiv221010352D/abstract"}, {"title": "Exploring Vision-Language Models for Open-Vocabulary Zero-Shot Action Segmentation", "url": "https://arxiv.org/pdf/2602.21406"}]

I found the paper you're looking for. "Temporal Action Segmentation: An Analysis of Modern Techniques" by Guodong Ding, Fadime Sener, and Angela Yao aims at densely identifying video frames in minutes-long videos with multiple action classes.

Despite the rapid growth of TAS techniques in recent years, no systematic survey had been conducted in these sectors, so this survey analyzes and summarizes the most significant contributions and trends. The survey examines the task definition, common benchmarks, types of supervision, and prevalent evaluation measures, and systematically investigates two essential techniques: frame representation and temporal modeling.

The paper is available on arXiv at <https://arxiv.org/abs/2210.10352> and was later published in IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI) in 2023.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. assistant (2026-06-13T20:19:58.219Z)

Both verified. The Ding et al. survey and MS-TCN are real with correct arXiv IDs, authors, venues. Let me verify the next batch: ASRF, Activity Grammars, and GAP.

7. user (2026-06-13T20:21:13.203Z)

Web search results for query: "Alleviating Over-segmentation Errors by Detecting Action Boundaries ASRF Ishikawa Kasai Aoki Kataoka WACV 2021 arXiv 2007.06866"

Links: [{"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries", "url": "https://yiskw713.github.io/asrf/"}, {"title": "WACV 2021 Open Access Repository", "url": "https://openaccess.thecvf.com/content/WACV2021/html/Ishikawa_Alleviating_Over-Segmentation_Errors_by_Detecting_Action_Boundaries_WACV_2021_paper.html"}, {"title": "Alleviating over-segmentation errors by detecting action boundaries ? Keio University", "url": "https://keio.elsevierpure.com/en/publications/alleviating-over-segmentation-errors-by-detecting-action-boundari-2"}, {"title": "GitHub - yiskw713/asrf: Code for 'Alleviating Over-segmentation Errors by Detecting Action Boundaries' accepted in WACV2021", "url": "https://github.com/yiskw713/asrf"}, {"title": "[PDF] Alleviating Over-segmentation Errors by Detecting Action Boundaries | Semantic Scholar", "url": "https://www.semanticscholar.org/paper/Alleviating-Over-segmentation-Errors-by-Detecting-Ishikawa-Kasai/4fed0235614df46fa0c53e05b71e9939b53a1bdf"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries", "url": "https://openaccess.thecvf.com/content/WACV2021/papers/Ishikawa_Alleviating_Over-Segmentation_Errors_by_Detecting_Action_Boundaries_WACV_2021_paper.pdf"}, {"title": "Alleviating Over-Segmentation Errors by Detecting Action Boundaries", "url": "https://papertalk.org/papertalks/27315"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries | Request PDF", "url": "https://www.researchgate.net/publication/352386005_Alleviating_Over-segmentation_Errors_by_Detecting_Action_Boundaries"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries", "url": "https://arxiv.org/pdf/2007.06866"}]

I found the paper you're looking for! Here's a summary of the research:

API Error: The socket connection was closed unexpectedly. For more information, pass `verbose: true` in the second argument to fetch()

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

8. user (2026-06-13T20:21:14.231Z)

Web search results for query: "Activity Grammars for Temporal Action Segmentation Gong Lee Jung Kwak Cho NeurIPS 2023 arXiv 2312.04266"

Links: [{"title": "Activity Grammars for Temporal Action Segmentation", "url": "https://jslee0424.github.io/KARI/"}, {"title": "[2312.04266] Activity Grammars for Temporal Action Segmentation", "url": "https://arxiv.org/abs/2312.04266"}, {"title": "NeurIPS Poster Activity Grammars for Temporal Action Segmentation", "url": "https://neurips.cc/virtual/2023/poster/70459"}, {"title": "Activity Grammars for Temporal Action Segmentation", "url": "https://proceedings.neurips.cc/paper_files/paper/2023/hash/ee6c4b99b4c0d3d60efd22c1ecdd9891-Abstract-Conference.html"}, {"title": "Activity Grammars for Temporal Action Segmentation", "url": "https://arxiv.org/pdf/2312.04266"}, {"title": "Activity Grammars for Temporal Action Segmentation - Kyungpook National University(KNU)", "url": "https://research.knu.ac.kr/en/publications/activity-grammars-for-temporal-action-segmentation"}, {"title": "Towards Generalizing Temporal Action Segmentation to Unseen Views", "url": "https://arxiv.org/html/2504.02512v1"}, {"title": "GitHub - gongda0e/KARI: Activity Grammars for Temporal Action Segmentation (NeurIPS 2023) - GitHub", "url": "https://github.com/gongda0e/kari"}, {"title": "Improving Temporal Action"}]

Segmentation and Detection with Hierarchical Task Grammar |
SpringerLink", "url": "https://link.springer.com/chapter/10.1007/978-3-031-78110-0_13"]}]

API Error: The socket connection was closed unexpectedly. For more information, pass `verbose: true` in the second argument to fetch()

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

9. assistant (2026-06-13T20:21:18.952Z)

Both verified. ASRF (arXiv:2007.06866, WACV 2021, Ishikawa/Kasai/Aoki/Kataoka ? official site and GitHub confirm) and Activity Grammars (arXiv:2312.04266, NeurIPS 2023, Gong et al. ? note GitHub repo "KARI" confirms authors Gong Dayoung et al.) are real. Let me verify GAP, BMN, and the SoccerNet action spotting paper.

10. user (2026-06-13T20:21:26.837Z)

Web search results for query: "Post-Processing Temporal Action Detection GAP Gaussian Approximated Nag Zhu Song Xiang CVPR 2023 arXiv 2211.14924"

Links: [{"title": "CVPR 2023 Open Access Repository", "url": "https://openaccess.thecvf.com/content/CVPR2023/html/Nag_Post-Processing_Temporal_Action_Detection_CVPR_2023_paper.html"}, {"title": "GAP: Post-Processing Temporal Action Detection", "url": "https://sauradip.github.io/project_pages/GAP/"}, {"title": "GitHub - sauradip/GAP: [CVPR 2023] Official PyTorch implementation of the paper \"GAP: Post-Processing Temporal Action Detection\"", "url": "https://github.com/sauradip/GAP"}, {"title": "[2211.14924] Post-Processing Temporal Action Detection", "url": "https://arxiv.org/abs/2211.14924"}, {"title": "Post-Processing Temporal Action Detection Sauradip Nag^{1,2} Xiatian Zhu^{1,3}", "url": "https://openaccess.thecvf.com/content/CVPR2023/papers/Nag_Post-Processing_Temporal_Action_Detection_CVPR_2023_paper.pdf"}, {"title": "Post-Processing Temporal Action Detection", "url": "https://arxiv.org/pdf/2211.14924"}, {"title": "TemporalMaxer: Maximize Temporal Context with only Max Pooling for Temporal Action Localization", "url": "https://arxiv.org/pdf/2303.09055"}, {"title": "Supplementary: Post-Processing Temporal Action Detection Sauradip Nag^{1,2}", "url": "https://openaccess.thecvf.com/content/CVPR2023/supplemental/Nag_Post-Processing_Temporal_Action_CVPR_2023_supplemental.pdf"}]

I found the paper you're looking for! Here's the information:

****Paper Details:****

The paper is titled "Post-Processing Temporal Action Detection" by Sauradip Nag, Xiatian Zhu, Yi-Zhe Song, and Tao Xiang, published in the Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) 2023. The arXiv preprint ID is 2211.14924.

****Main Contribution:****

The paper addresses the problem of temporal quantization error introduced when existing Temporal Action Detection (TAD) methods convert varying-length videos into fixed-length snippet representation sequences, which temporally downsamples the video and reduces inference resolution.

The authors introduce a novel model-agnostic post-processing method and an efficient Taylor-expansion based approximation, dubbed as Gaussian Approximated Post-processing (GAP).

****Results:****

Experiments demonstrate that GAP can consistently improve a wide variety of pre-trained off-the-shelf TAD models on ActivityNet (+0.2% 0.7% in average mAP) and THUMOS (+0.2% 0.5% in average mAP) benchmarks, with performance gains that are already significant and highly comparable to those achieved by novel model designs.

The code is available at <https://github.com/sauradip/GAP>.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

11. user (2026-06-13T20:21:27.124Z)

Web search results for query: "BMN Boundary-Matching Network Temporal Action Proposal Generation Lin Liu Li Ding Wen ICCV 2019 arXiv 1907.09702"

Links: [{"title": "[1907.09702] BMN: Boundary-Matching Network for Temporal Action Proposal Generation", "url": "https://arxiv.org/abs/1907.09702"}, {"title": "BMN: Boundary-Matching Network for Temporal Action Proposal Generation | IEEE Conference Publication | IEEE Xplore", "url": "https://ieeexplore.ieee.org/iel7/8972782/9008105/09009808.pdf"}, {"title": "BMN: Boundary-Matching Network for Temporal Action Proposal Generation | Request PDF", "url": "https://www.researchgate.net/publication/334644585_BMN_Boundary-Matching_Network_for_Temporal_Action_Proposal_Generation"}, {"title": "BMN: Boundary-Matching Network for Temporal Action Proposal Generation", "url": "https://arxiv.org/pdf/1907.09702"}, {"title": "BMN: Boundary-Matching Network for Temporal Action Proposal Generation", "url": "https://openaccess.thecvf.com/content_ICCV_2019/papers/Lin_BMN_Boundary-Matching_Network_for_Temporal_Action_Proposal_Generation_ICCV_2019_paper.pdf"}, {"title": "[Quick Review] BMN: Boundary-Matching Network for Temporal Action Proposal Generation", "url": "https://liner.com/review/bmn-boundarymatching-network-for-temporal-action-proposal-generation"}, {"title": "GitHub - xlliu7/BMN-Boundary-Matching-Network: A Fast PyTorch implementation for ICCV 19 paper \"BMN: Boundary-Matching Network for Temporal Action Proposal Generation\"", "url": "https://github.com/xlliu7/BMN-Boundary-Matching-Network"}, {"title": "Fast Learning of Temporal Action Proposal via Dense Boundary Generator", "url": "https://arxiv.org/pdf/1911.04127"}]

I found the paper you're looking for. Here are the key details:

Paper Title: BMN: Boundary-Matching Network for Temporal Action Proposal Generation

Authors: Tianwei Lin, Xiao Liu, Xin Li, Errui Ding, and Wen

Publication: ICCV 2019

arXiv ID: 1907.09702

Main Contribution: The paper introduces the Boundary-Matching (BM) mechanism to evaluate confidence scores of densely distributed proposals, which denote a proposal as a matching pair of starting and ending boundaries and combine all densely distributed BM pairs into the BM confidence map.

Research Focus: Temporal action proposal generation is a challenging and promising task which aims to locate temporal regions in real-world videos where action or event may occur.

Results: BMN was evaluated on two challenging datasets: THUMOS-14 and ActivityNet-1.3, where BMN shows significant performance improvement with remarkable efficiency and generalizability.

The paper is available on arXiv at <https://arxiv.org/abs/1907.09702> and the ICCV proceedings.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

12. assistant (2026-06-13T20:21:32.264Z)

Both verified. GAP (arXiv:2211.14924, CVPR 2023, Nag/Zhu/Song/Xiang ? confirmed) and BMN (arXiv:1907.09702, ICCV 2019, Lin et al. ? confirmed; note the survey omits "Wen Sun" but the authors list and id are correct). Let me now verify the SoccerNet action spotting paper and the two badminton-specific papers, which are the highest fabrication-risk items (auth-gated DOIs).

13. user (2026-06-13T20:21:40.776Z)

Web search results for query: "Service and End of Rally Detection in Badminton" Springer 2024 10.1007/978-3-031-69073-0_15"

Links: [{"title": "Service and End of Rally Detection in Badminton Videos | Springer Nature Link", "url": "https://link.springer.com/chapter/10.1007/978-3-031-69073-0_15"}, {"title": "A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton Videos - Monash University", "url": "https://research.monash.edu/en/publications/a-comparative-analysis-of-deep-learning-models-and-gradient-compu"}, {"title": "A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton Videos | SN Computer Science | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s42979-025-03935-0"}, {"title": "Detection of service activity in a badminton game | Request PDF", "url": "https://www.researchgate.net/publication/261087356_Detection_of_service_activity_in_a_badminton_game"}, {"title": "Badminton Detection Using Lightweight Neural Networks for Service Fault Judgement | SpringerLink", "url": "https://link.springer.com/chapter/10.1007/978-981-97-2275-4_15"}, {"title": "A Machine Learning Framework for Shuttlecock Tracking and Player Service Fault Detection | Springer Nature Link", "url": "https://link.springer.com/chapter/10.1007/978-3-031-39059-3_5"}, {"title": "Estimation of control area in badminton doubles with pose information from top and back view drone videos | Multimedia Tools and Applications | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s11042-023-16362-1"}]

I found the article you're looking for! Here are the details:

The paper "Service and End of Rally Detection in Badminton Videos" by See, S.Y., Hassan, M.F., Krishnasamy, G., Loo, J.Y., and Paramesran, R. was published in 2024 in the ISACE 2024 conference proceedings as part of Lecture Notes in Computer Science, vol 14794, by Springer, Cham, with DOI 10.1007/978-3-031-69073-0_15.

****Key findings from the paper:****

Gradients are used to detect the flight of the shuttlecock till it eventually comes to rest which signifies the end of the rally. In computing the end of the rally, the algorithm considers the missing shuttlecock flight coordinates and when the shuttlecock is out of camera view.

Two videos of 4 and 12 minutes comprising 12 and 23 consecutive points were used to determine the accuracy of the proposed algorithm, and the experimental results show 88.6% accuracy for a total of 35 rally points.

The article is available on SpringerLink at the URL you provided.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

14. user (2026-06-13T20:21:43.202Z)

Web search results for query: "Action Spotting using Dense Detection Anchors Revisited Soares Shah Biswas 2022 arXiv 2206.07846 SoccerNet"

Links: [{"title": "ASTRA: An Action Spotting TRAnsformer for Soccer Videos", "url": "https://arxiv.org/html/2404.01891"}, {"title": "[PDF] Action Spotting using Dense Detection Anchors Revisited: Submission to the SoccerNet Challenge 2022 | Semantic Scholar", "url": "https://www.semanticscholar.org/paper/Action-Spotting-using-Dense-Detection-Anchors-to-Soares-Shah/196e13d9ada026058e9be6a8a4822ec998327ba5"}, {"title": "Deep learning for action spotting in association football videos", "url": "https://arxiv.org/pdf/2410.01304"}, {"title": "[2206.07846] Action Spotting using Dense Detection Anchors Revisited: Submission to the SoccerNet Challenge 2022", "url": "https://arxiv.org/abs/2206.07846"}, {"title": "Action Spotting using Dense Detection Anchors Revisited: Submission to the SoccerNet Challenge 2022", "url": "https://explore.openaire.eu/search/publication?pid=10.48550/arxiv.2206.07846"}, {"title": "Action Spotting using Dense Detection Anchors Revisited: Submission to the SoccerNet Challenge 2022", "url": "https://arxiv.org/pdf/2206.07846"}, {"title": "[2205.10450v1] Temporally Precise"}]

Action Spotting in Soccer Videos Using Dense Detection Anchors", "url": "http://export.arxiv.org/abs/2205.10450v1"}, {"title": "[2404.01891] ASTRA: An Action Spotting TRAnsformer for Soccer Videos", "url": "https://arxiv.org/html/2404.01891"}, {"title": "GitHub - yahoo/spivak · GitHub", "url": "https://github.com/yahoo/spivak"}]

I found the paper you're looking for. Here's what the search results show:

The paper "Action Spotting using Dense Detection Anchors Revisited: Submission to the SoccerNet Challenge 2022" was written by João V. B. Soares and Avijit Shah from Yahoo Research and describes their submission to the Action Spotting SoccerNet Challenge 2022, which was part of the CVPR 2022 ActivityNet Workshop.

The submission is based on a method that focuses on increasing temporal precision via a densely sampled set of detection anchors, and due to its emphasis on temporal precision, this approach showed significant improvements in the tight average-mAP metric.

The changes they introduced helped them achieve first place in the challenge and led to a new state-of-the-art on SoccerNet's test set when using the dataset's standard experimental protocol.

Note: The search results also show there's a related paper by Soares, Shah, and Topojoy Biswas that presents a model for temporally precise action spotting in videos using a dense set of detection anchors, predicting detection confidence and corresponding fine-grained temporal displacement for each anchor.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

15. assistant (2026-06-13T20:21:48.593Z)

Both verified as real. Important detail on the SoccerNet paper: the survey's title says "Action Spotting using Dense Detection Anchors Revisited" ? the real full title adds ": Submission to the SoccerNet Challenge 2022", authors confirmed as Soares, Shah, Biswas, arXiv:2206.07846. The badminton Springer 2024 paper (DOI 10.1007/978-3-031-69073-0_15) is real, authors See et al.

Now the last cited paper ? the SN Computer Science 2025 badminton paper ? and the one inline arXiv reference in the pitfalls (2106.14118, audio-augmented TAL fusion).

16. user (2026-06-13T20:21:56.331Z)

Web search results for query: "arXiv 2106.14118 audio visual temporal action localization fusion"

Links: [{"title": "Multi-Resolution Audio-Visual Feature Fusion for Temporal Action Localization", "url": "https://arxiv.org/pdf/2310.03456"}, {"title": "[2106.14118v2] Hear Me Out: Fusional Approaches for Audio Augmented Temporal Action Localization", "url": "https://export.arxiv.org/abs/2106.14118v2"}, {"title": "[2106.14118] Hear Me Out: Fusional Approaches for Audio Augmented Temporal Action Localization", "url": "https://arxiv.org/abs/2106.14118"}, {"title": "[2310.03456] Multi-Resolution Audio-Visual Feature Fusion for Temporal Action Localization", "url": "https://arxiv.org/abs/2310.03456"}, {"title": "(PDF) Multi-Resolution Audio-Visual Feature Fusion for Temporal Action Localization", "url": "https://www.academia.edu/12446179/8/Multi_Resolution_Audio_Visual_Feature_Fusion_for_Temporal_Action_Localization"}, {"title": "UniA V: Unified Audio-Visual Perception for Multi-Task Video Event Localization", "url": "https://arxiv.org/pdf/2404.03179"}, {"title": "Multi-Resolution Audio-Visual Feature Fusion for Temporal Action Localization", "url": "https://andrewjohngilbert.github.io/Multi-TAL/"}, {"title": "TALLFormer: Temporal Action Localization with a Long-memory Transformer", "url": "https://arxiv.org/pdf/2204.01680"}, {"title": "Hear Me Out: Fusional Approaches for Audio Augmented Temporal Action Localization", "url": "https://arxiv.org/pdf/2106.14118"}]

Based on the search results, I found the paper you're looking for:

The paper arXiv:2106.14118 is titled "Hear Me Out: Fusional Approaches for Audio Augmented Temporal Action Localization" by Anurag Bagchi, Jazib Mahmood, Dolton Fernandes, and Ravi Kiran Sarvadevabhatla.

****Key details about the paper:****

The paper addresses temporal action localization (TAL) in untrimmed videos by incorporating the audio modality, which previous state-of-the-art approaches had left unexploited while only considering RGB and Flow modalities.

To the best of the authors' knowledge, their work is the first to jointly consider audio and video modalities for supervised TAL. Their approaches achieve state-of-the-art performance on large-scale benchmark datasets - ActivityNet-1.3 (54.34 mAP@0.5) and THUMOS14 (57.18 mAP@0.5).

The experiments include ablations involving multiple fusion schemes, modality combinations and TAL architectures.

The paper was submitted on June 27, 2021 and last revised on October 17, 2021.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

17. user (2026-06-13T20:21:57.817Z)

Web search results for query: "'A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton" SN Computer Science 10.1007/s42979-025-03935-0"

Links: [{"title": "A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton Videos - Monash University", "url": "https://research.monash.edu/en/publications/a-comparative-analysis-of-deep-learning-models-and-gradient-compu"}, {"title": "A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton Videos | SN Computer Science | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s42979-025-03935-0"}, {"title": "Comparative analysis of automated foul detection in football using deep learning architectures", "url": "https://www.ncbi.nlm.nih.gov/pmc/articles/PMC12022285/"}, {"title": "An AI-based badminton smash detection using residual-shuffle network optimized based on upgraded pufferfish optimizer - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/pii/S2090447925001558"}, {"title": "High Speed and Tiny Objects Tracking System in Racquet Sports Videos Using Deep Learning with Trajectory Rectification Feature - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/pii/S1877050925016709"}, {"title": "Automated Service Height Fault Detection Using Computer Vision and Machine Learning for Badminton Matches", "url": "https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10747833/"}, {"title": "Strategy analysis of badminton players using deep learning from IMU and UWB wearables - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/abs/pii/S2542660524002014"}]

I found the research paper you're looking for! Here's a summary of "A Comparative Analysis of Deep Learning Models and Gradient Computation for Rally Detection in Badminton":

Overview

Badminton videos are commonly used by coaches and players to analyze performance and gain valuable insights into each point played in a match, though it is time consuming to manually annotate the start of service and end of each rally.

Methodology

The research used two different models, Bidirectional Long Short Term Memory (BiLSTM) and Convolutional Neural Network (CNN), for detecting the start of service while gradient computation determined the end of the rally. A total of 12 frames with 26 features from each frame comprising the distance, angle, gradient and movements that are typical of serving and receiving actions was arranged in a 2D format that was used as inputs to the models.

Results

The experimental results demonstrated the highest classification accuracy of approximately 95%

for the BiLSTM model, with frame intervals of 4, while the CNN achieved similar accuracy at frame intervals of 5. For the detection of the end of a rally, at least 80% of consecutive frames needed to exhibit no motion of the shuttlecock, achieving an average accuracy of 95% across a total of 187 rallies in the test dataset.

****Citation:**** SN Computer Science, 6(5), Article 447. <https://doi.org/10.1007/s42979-025-03935-0>

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

18. assistant (2026-06-13T20:22:05.331Z)

All 12 cited works (10 key_papers + 2 inline arXiv references) are verified real with correct IDs, authors, and venues. No fabrications. Notably, the SN Computer Science 2025 paper's claims (BiLSTM service detection ~95%, end-of-rally via $\geq 80\%$ consecutive no-motion frames, 187 rallies) all check out exactly as described.

Now the harder adversarial task: refuting the "partial-overlap" verdict, with focus on the survey's claim that the specific assembly (physics cues as feature columns + tiny logistic-regression scorer + reconciler + LOVO harness for scarce-label badminton rally segmentation) is a "novel-combination." Let me hunt for closer prior art on rule-based/grammar/constraint post-processing correctors AND badminton/racket-sports rally segmentation systems that may already combine these elements.

19. user (2026-06-13T20:22:15.260Z)

Web search results for query: "tennis broadcast play break segmentation rules domain knowledge post-processing correction shuttle ball inactivity"

Links: [{"title": "AUDIOVISUAL INTEGRATION FOR TENNIS BROADCAST STRUCTURING", "url": "https://www.irisa.fr/metiss/publis/kijak-cbmi03.pdf"}, {"title": "Method and system for high-level structure analysis and event detection in domain specific videos", "url": "https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/6813313"}, {"title": "Learning Wheelchair Tennis Navigation from Broadcast Videos with Domain Knowledge Transfer and Diffusion Motion Planning", "url": "https://arxiv.org/html/2409.19771v2"}, {"title": "TennisExpert: Towards Expert-Level Analytical Sports Video Understanding", "url": "https://arxiv.org/pdf/2603.13397"}, {"title": "(PDF) Automated Ball Tracking in Tennis Videos", "url": "https://www.academia.edu/33934086/Automated_Ball_Tracking_in_Tennis_Videos"}, {"title": "(PDF) Analyzing Racket Sports From Broadcast Videos", "url": "https://www.academia.edu/70243184/Analyzing_Racket_Sports_From_Broadcast_Videos"}, {"title": "Beyond the Broadcast: Enhancing VR Tennis Broadcasting through Embedded Visualizations and Camera Techniques", "url": "https://arxiv.org/pdf/2507.20006"}, {"title": "Scene segment dividing device", "url": "https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/8306390"}, {"title": "(PDF) Players tracking and ball detection for an automatic tennis video annotation", "url": "https://www.researchgate.net/publication/221144146_Players_tracking_and_ball_detection_for_an_automatic_tennis_video_annotation"}]

I found several relevant research papers and patents on tennis broadcast video analysis with play-break segmentation and domain knowledge. Here are the key findings:

Play-Break Segmentation Rules

A break happens at the end of at least ten consecutive points. Considering these rules in the model ensures a long-time correlation between shots, and avoids, for example, the apparition of an interleaving of points and breaks.

Detection Methods

Play-break segmentation is performed based on different views, with close-up view typically corresponding to a break. Additionally, a lightweight, purely audio-based detection model is employed, with a neural network trained to identify the distinct acoustic signatures of racquet-ball impacts from the broadcast audio track.

Ball Inactivity and Segmentation

Ball movement vector of ball track information is used to extract the ball fixed-time stop point, thereby specifying the break segment. Once hitting moments are identified, these temporal anchors are clustered in the temporal domain to form candidate rally intervals, and specific filtering rules are then applied to retain valid segments corresponding to complete, uninterrupted rallies.

Post-Processing Correction

Minor perturbations and occasional missing player data (due to occlusions) are corrected using k-nearest neighbors interpolation ($k=5$), followed by smoothing with a 5-frame moving average filter.

Domain Knowledge Application

Sports video processing is an interesting topic for research, since the clearly defined game rules in sports provide the rich domain knowledge for analysis. A score-based indexing system segments tennis matches into rallies and recognizes scores using contextual knowledge of the tennis scoring system.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

20. user (2026-06-13T20:22:16.794Z)

Web search results for query: "badminton rally segmentation shuttlecock trajectory player audio multi-cue feature fusion classifier serve detection"

Links: [{"title": "Future Prediction of Shuttlecock Trajectory in Badminton Using Player?s Information", "url": "https://www.mdpi.com/2313-433X/9/5/99"}, {"title": "Enhancing Badminton Game Analysis: An Approach to Shot Refinement via a Fusion of Shuttlecock Tracking and Hit Detection from Monocular Camera - PMC", "url": "https://pmc.ncbi.nlm.nih.gov/articles/PMC11244353/"}, {"title": "Enhancing Badminton Game Analysis: An Approach to Shot Refinement via a Fusion of Shuttlecock Tracking and Hit Detection from Monocular Camera", "url": "https://www.mdpi.com/1424-8220/24/13/4372"}, {"title": "(PDF) Enhancing Badminton Game Analysis: An Approach to Shot Refinement via a Fusion of Shuttlecock Tracking and Hit Detection from Monocular Camera", "url": "https://www.researchgate.net/publication/382039345_Enhancing_Badminton_Game_Analysis_An_Approach_to_Shot_Refinement_via_a_Fusion_of_Shuttlecock_Tracking_and_Hit_Detection_from_Monocular_Camera"}, {"title": "Enhancing Badminton Game Analysis: An Approach to Shot Refinement via a Fusion of Shuttlecock Tracking and Hit Detection from Monocular Camera", "url": "https://www.ncbi.nlm.nih.gov/pmc/articles/PMC11244353/"}, {"title": "TrackNetV3: Enhancing Shuttlecock Tracking with Augmentations and Trajectory Rectification | Proceedings of the 5th ACM International Conference on Multimedia in Asia", "url": "https://dl.acm.org/doi/abs/10.1145/3595916.3626370"}]

I found relevant search results on badminton video analysis with focus on rally segmentation and shuttlecock tracking. Here's what the search returned:

Key Findings:

****Advanced Methods for Badminton Analysis:****

Recent research combines TrackNet (a deep neural network for tracking small objects) to extract shuttlecock flight trajectory with the YOLOv7 model to identify player swinging, proposing a shot refinement algorithm to obtain the correct hitting moment.

****Performance Results:****

A recent study achieved an accuracy of 89.7%, a recall rate of 91.3%, and an F1 rate of 90.5% in

69 matches with 1582 rallies of Badminton World Federation (BWF) match videos. Shot type classification achieved accuracy rates of 72.1%, 65.4%, and 54.1% at different thresholds.

****Related Capabilities:****

Systems like Coach AI utilize deep learning methodologies to extract player movements and positioning while detecting shuttlecock trajectory for identifying swing moments, rally segmentation, and shot categorization.

****Future Research Needs:****

Automating rally segmentation and shuttlecock position detection in match videos remains an important area for future work.

The search results focus primarily on vision-based approaches combining shuttlecock tracking and player motion detection, though the results don't specifically discuss audio-based features or multi-cue feature fusion with audio data as mentioned in your original query.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

21. assistant (2026-06-13T20:22:22.412Z)

Interesting prior art surfacing. The Kijak et al. tennis broadcast structuring paper (IRISA) uses audiovisual integration with play/break (rally) structure rules ? a strong analog. And there's broader "domain rules + ball inactivity + candidate clustering + filtering rules" prior art in racket sports. Let me dig into two more refutation angles: (1) the specific "rule-based / constraint post-processing corrector on a base segmenter" framing in TAS beyond Activity Grammars, and (2) whether a published badminton/tennis system already combines multi-cue features + a lightweight learned scorer + leave-one-out on scarce labels.

22. user (2026-06-13T20:22:30.723Z)

Web search results for query: "badminton tennis rally detection shuttlecock no-motion play break leave-one-video-out cross validation small dataset logistic regression interpretable features"

Links: [{"title": "Future Prediction of Shuttlecock Trajectory in Badminton Using Player's Information", "url": "https://www.mdpi.com/2313-433X/9/5/99"}, {"title": "Automatic rally detection and scoring", "url": "https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/10751601"}, {"title": "Badminton shuttlecock detection and prediction of trajectory using multiple 2 dimensional scanners | Request PDF", "url": "https://www.researchgate.net/publication/301710900_Badminton_shuttlecock_detection_and_prediction_of_trajectory_using_multiple_2_dimensional_scanners"}, {"title": "Enhancing Badminton Game Analysis: An Approach to Shot Refinement via a Fusion of Shuttlecock Tracking and Hit Detection from Monocular Camera", "url": "https://www.mdpi.com/1424-8220/24/13/4372"}, {"title": "MonoTrack: Shuttle trajectory reconstruction from monocular badminton video | Request PDF", "url": "https://www.researchgate.net/publication/362900461_MonoTrack_Shuttle_trajectory_reconstruction_from_monocular_badminton_video"}, {"title": "One-Shot Badminton Shuttle Detection for Mobile Robots", "url": "https://arxiv.org/pdf/2603.06691"}, {"title": "Enhancing Badminton Game Analysis: An Approach to Shot Refinement via a Fusion of Shuttlecock Tracking and Hit Detection from Monocular Camera - PMC", "url": "https://pmc.ncbi.nlm.nih.gov/articles/PMC11244353/"}]

The search results show research related to badminton video analysis, shuttle detection, and rally identification, but don't specifically match all the particular terms in your query (especially leave-one-video-out cross validation, logistic regression for interpretable features, and detection of no-motion play breaks).

The results indicate that current research in this area uses deep learning approaches like TrackNet, YOLO, and other neural networks for shuttlecock tracking and hit detection. For example, one system uses TrackNet for tracking the shuttlecock and YOLOv7 to identify player swinging, while another method achieved 89.7% accuracy and 91.3% recall on badminton match videos.

However, the specific combination of techniques you're asking about? particularly logistic

regression for interpretable features with leave-one-video-out cross-validation on a small dataset for detecting no-motion play breaks?doesn't appear in these results. This might be a specialized research paper or methodology that isn't widely indexed, or it may be a very recent or niche application. You may want to try searching directly on academic databases like arXiv, Google Scholar, or specific conference proceedings if you're looking for a particular paper.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

23. user (2026-06-13T20:22:33.959Z)

Web search results for query: "rule-based post-processing correction over-segmentation temporal action segmentation constraint satisfaction prior knowledge inference time refine base model"

Links: [{"title": "Improving Temporal Action Segmentation via Constraint-Aware Decoding", "url": "https://arxiv.org/html/2605.10149"}, {"title": "NeurIPS Poster OnlineTAS: An Online Baseline for Temporal Action Segmentation", "url": "https://neurips.cc/virtual/2024/poster/94466"}, {"title": "OnlineTAS: An Online Baseline for Temporal Action Segmentation", "url": "https://arxiv.org/pdf/2411.01122"}, {"title": "DistNet2D: Leveraging long-range temporal information for efficient segmentation and tracking", "url": "https://arxiv.org/pdf/2310.19641"}, {"title": "Maximization and restoration: Action segmentation through dilation passing and temporal reconstruction - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/abs/pii/S003132032200245X"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries", "url": "https://arxiv.org/pdf/2007.06866"}, {"title": "A video action segmentation method based on multi-scale temporal features - IOPscience", "url": "https://iopscience.iop.org/article/10.1088/2631-8695/ae68d8"}, {"title": "Temporal Action Segmentation: An Analysis of Modern Techniques", "url": "https://arxiv.org/pdf/2210.10352"}, {"title": "Improving action segmentation via explicit similarity measurement - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/abs/pii/S0952197625032828"}, {"title": "Temporal Action Segmentation: A Comprehensive Overview", "url": "https://www.emergentmind.com/topics/temporal-action-segmentation"}]

Based on the web search results, I found relevant research on this topic. Here are the key findings:

Rule-Based Post-Processing for Temporal Action Segmentation

There is a rule-based refinement framework that offers a decoding strategy to inject structural knowledge into temporal action segmentation (TAS) predictions using data-derived constraints, applied as a post-processing step on top of existing TAS models without requiring model modification or retraining.

Addressing Over-Segmentation

Post-processing methods mitigate over-segmentation by adjusting action boundaries according to network predictions and action length. Additionally, restoration approaches focus on correcting inconsistent labels in order to reduce over-segmentation.

Constraint-Based Approaches

Some methods introduce temporal logic loss that measures the consistency between the output and imposed temporal constraints. Recent approaches use position, boundary, and relation priors of human actions to enhance segmentation results.

Inference-Time Refinement

Explicit boundary detection/regression and post-processing techniques suppress micro-segmentation and sharpen boundaries at inference time.

The most recent relevant work appears to be from a May 2026 paper on "Improving Temporal Action Segmentation via Constraint-Aware Decoding," which directly addresses this topic of using

constraints and rule-based post-processing to improve temporal action segmentation models.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

24. assistant (2026-06-13T20:22:40.350Z)

Two significant finds that bear on the verdict:

1. **"Improving Temporal Action Segmentation via Constraint-Aware Decoding"** (arXiv:2605.10149, May 2026) ? this is a NEWER and even closer analog to the reconciler than Activity Grammars: a rule-based refinement framework injecting structural knowledge as post-processing on top of any TAS model without retraining, using data-derived constraints. This is essentially the reconciler concept, published. The survey didn't cite it (its own search query mentions "constraint-aware decoding" but it's not in key_papers).

2. **"Automatic rally detection and scoring"** (US patent 10751601) ? a patent directly on rally detection/scoring. Worth checking whether it combines cues the way the survey claims is novel.

Let me verify the constraint-aware decoding paper is real and characterize it, and check the rally-detection patent.

25. user (2026-06-13T20:22:45.779Z)

Paper Summary

Title: "Improving Temporal Action Segmentation via Constraint-Aware Decoding"

Authors: Yeo Keat Ee, Debaditya Roy, Chen Li, Hao Zhang, Basura Fernando

Year: 2026 (submitted May 11, 2026)

Venue: ICPR 2026

Method

Yes, this approach applies constraint-based post-processing without model retraining. The paper describes "inference-time refinement without retraining or added model complexity."

The method resolves constraints **globally** by integrating them into "a modified Viterbi decoding algorithm," which is a dynamic programming approach that finds optimal solutions across the entire sequence.

Specific Constraints

The framework uses three statistical structural priors extracted directly from annotated data:

- Transition confidence** ? likelihood of action transitions
- Action boundary sets** ? permitted boundaries between actions
- Per-class duration** ? typical length constraints for each action class

26. user (2026-06-13T20:22:48.979Z)

Web search results for query: "Automatic rally detection and scoring" patent 10751601 method cues shuttle player audio"

Links: [{"title": "Automatic rally detection and scoring", "url": "https://image-pubs.uspto.gov/dirsearch-public/print/downloadPdf/10751601"}, {"title": "US20180117440A1 - Automatic rally detection and scoring - Google Patents", "url": "https://patents.google.com/patent/US20180117440A1/en"}, {"title": "Automatic rally detection on broadcast tennis videos | Request PDF", "url": "https://www.researchgate.net/publication/290310875_Automatic_rally_detection_on_broadcast_tennis_videos"}, {"title": "A New Perspective"}]

for Shuttlecock Hitting Event

Detection", "url": "https://arxiv.org/pdf/2306.10293"}, {"title": "Table Tennis Patents and Patent Applications (Class 473/475) - Justia Patents Search", "url": "https://patents.justia.com/patents-by-us-classification/473/475"}]

Based on the search results, I found information about patent 10751601 on "Automatic Rally Detection and Scoring."

The patent describes a computer-implemented method for detecting a rally in a sports game that involves wirelessly receiving first motion data detected by a sensor attached to a sports instrument, classifying each detected stroke action using a trained stroke classification model, and determining that certain detected non-stroke actions are intentional special user actions according to customized definitions. The method detects a rally using a stroke buffer at the client device based on classified stroke actions and intentional special user actions, where the rally comprises a sequence of one or more strokes detected during the sports game.

The method also includes responsive mechanisms for detecting the rally: when strokes are detected, the stroke buffer status is checked; when full, the first stroke is deleted; a new stroke is stored; and a trained machine learning model is applied to stroke actions in the buffer to determine a rally state indicative of the rally.

However, the search results I found do not contain specific information about the audio cues for the shuttle player or specific scoring methods mentioned in your query. The document excerpt provided focuses on the motion detection and rally detection methodology rather than the audio cueing features you're asking about.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

27. assistant (2026-06-13T20:22:57.799Z)

Good. The patent (10751601) is a wearable-sensor / IMU-based rally detector ? different modality (motion sensors on the racket, not static-camera video), so it is not close prior art to the video-segmentation architecture. I'll set it aside.

The constraint-aware decoding paper (arXiv:2605.10149, ICPR 2026, Yeo Keat Ee et al.) is confirmed real and is a closer analog to the reconciler than Activity Grammars in one specific way: it extracts structural priors (transition confidence, action boundary sets, per-class duration) directly from data and resolves them via a global Viterbi decode as post-processing on a frozen base model ? this matches the survey's own "techniques_to_adopt" recommendation to "frame reconciler rules as a grammar and resolve globally (Viterbi-style)." It is concurrent (May 2026), so it doesn't displace the survey's novelty claim, but it should be cited.

Let me do one more refutation pass on the most defensible novelty claim ? the "report-first, evidence-vs-decision contradiction" reconciler framing ? and check the Kijak tennis paper which looked like strong task+method prior art.

28. user (2026-06-13T20:23:04.773Z)

I apologize, but the provided content appears to be a PDF file in binary/compressed format that cannot be meaningfully parsed as plain text. The content shows PDF structure markers and encoded stream data rather than readable document text.

To answer your questions about the title, authors, year, segmentation methodology, audio-visual fusion, model type, and post-processing rules, I would need the document in a readable text format (such as plain text, HTML, or properly extracted PDF content).

Could you please provide the document in a more accessible format, or share the key textual content directly?

[Binary content (application/pdf, 315.1KB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\tool-results\webfetch-1781381820827-jqeefer.pdf]

29. user (2026-06-13T20:23:06.217Z)

Web search results for query: "sports video rally segmentation candidate proposals lightweight classifier handcrafted features scarce labels weakly supervised play break detection few matches"

Links: [{"title": "Unsupervised Temporal Feature Aggregation for Event Detection in Unstructured Sports Videos", "url": "https://arxiv.org/pdf/2002.08097"}, {"title": "Deep Learning for Sports Video Event Detection: Tasks, Datasets, Methods, and Challenges", "url": "https://arxiv.org/html/2505.03991v3"}, {"title": "TennisExpert: Towards Expert-Level Analytical Sports Video Understanding", "url": "https://arxiv.org/pdf/2603.13397"}, {"title": "A Survey on Content-Aware Video Analysis for Sports", "url": "https://arxiv.org/pdf/1703.01170"}, {"title": "(PDF) Contrastive Learning for Sports Video: Unsupervised Player Classification", "url": "https://www.researchgate.net/publication/351019512_Contrastive_Learning_for_Sports_Video_Unsupervised_Player_Classification"}, {"title": "Self-Learning for Player Localization in Sports Video", "url": "https://arxiv.org/pdf/1307.7198"}, {"title": "SmartTennisTV: Automatic indexing of tennis videos", "url": "https://arxiv.org/pdf/1801.01430"}, {"title": "A framework for flexible summarization of racquet sports video using multiple modalities - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/abs/pii/S1077314208001355"}, {"title": "Action Spotting and Precise Event Detection in Sports: Datasets, Methods, and Challenges", "url": "https://arxiv.org/html/2505.03991v1"}]

Based on the search results, your query appears to relate to research on sports video analysis, particularly focusing on rally segmentation and event detection in racquet sports. Here are the relevant findings:

Rally Segmentation in Sports Videos

Lightweight detection models have been employed to identify distinct acoustic signatures of racquet-ball impacts from broadcast audio tracks for efficient rally segmentation across large video volumes.

HOG (Histogram of Oriented Gradients) features? a handcrafted feature approach? have been used to train classifiers like γ^2 -kernel SVM to distinguish rally frames from non-rally frames in tennis broadcasts, with smoothing techniques applied to remove false positives/negatives.

Weakly Supervised and Limited Label Approaches

Self-supervised frameworks that exploit temporal consistency, contrastive objectives, or multimodal alignment can learn meaningful event representations from unlabeled or weakly labeled sports videos.

Self-learning approaches have proven particularly effective when very little labeled data is available, enabling fully automatic label acquisition from realistic broadcast sports footage with only small amounts of initial labeled data.

Play/Break Detection

Play-break segmentation has been used as a universal detection scope across sports, with racquet sports games being segmented into rally (play) and break events, with rallies then ranked according to their excitement levels.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

30. assistant (2026-06-13T20:23:15.858Z)

The SmartTennisTV (arXiv:1801.01430) result is a notable find ? automatic indexing of tennis videos, likely using rally segmentation with handcrafted features. And the search confirms a well-established sports-video pattern: HOG features + γ^2 -SVM rally/non-rally classification with temporal smoothing as post-processing, plus self-learning with little labeled data. This is the

"candidate windows + light classifier on handcrafted features + smoothing post-process for scarce-label sports rally segmentation" combination the survey calls a novel combination ? it has substantial precedent in racket sports.

I have enough to render a corrected verdict. Let me do one final targeted check on SmartTennisTV's method to confirm it as closer prior art for the overall assembly.

31. user (2026-06-13T20:23:22.436Z)

SmartTennisTV Paper Analysis

****Title, Authors, Year:****

"SmartTennisTV: Automatic indexing of tennis videos" by Anurag Ghosh and C.V. Jawahar (2018)

****Rally Segmentation Approach:****

The available abstract does not specify the technical mechanism for segmenting rallies. It states the system "temporally segments the rallies in the video and then recognizes the scores from each of the segments," but does not detail whether this uses shot/scene detection, learned classifiers, or other methods.

****Video Type & Post-Processing:****

The work targets broadcast tennis videos (BTVs). The abstract confirms rule-based refinement occurs: the approach applies "score refinement using the knowledge of the tennis scoring system" as post-processing to correct recognized scores.

****Limitation:****

The abstract alone cannot confirm whether shot/scene cues or learned classifiers drive the segmentation. The full paper (accessible via PDF link) would contain these implementation details, but that content is not provided in the excerpt shown.

32. assistant (2026-06-13T20:23:33.487Z)

I have sufficient evidence. SmartTennisTV is broadcast-video (not static-camera) and applies rule-based refinement using scoring-system knowledge as post-processing ? confirming the "rules as structural prior for post-correction" pattern in racket sports, though its segmentation mechanism (likely scene/shot-based on broadcast cuts) differs from this project's static-camera physics-cue approach.

Let me compile the verdict. Summary of findings:

****Fabrication check: clean.**** All 10 key_papers + 2 inline arXiv refs (2106.14118, and the implicitly-relevant 2605.10149) verified real with correct IDs/authors/venues. Two minor title imprecisions (SoccerNet paper's full title includes ": Submission to the SoccerNet Challenge 2022"; BMN author list drops "Wen Sun") ? not fabrications.

****Novelty refutation:**** The survey's verdict is already the conservative "partial-overlap," and it explicitly frames the specific scarce-label badminton instantiation as "novel-combination, not a clean gap." My searches mostly support keeping partial-overlap but surface closer prior art that the survey missed, which slightly strengthens the "reinventing" pull on two sub-claims (the reconciler, and the handcrafted-features + light-classifier + smoothing for scarce-label racket-sports rally segmentation). The overall verdict should remain partial-overlap.

33. user (2026-06-13T20:24:29.709Z)

Structured output provided successfully

34. assistant (2026-06-13T20:24:41.824Z)

Adversarial fact-check complete. Summary:

****Fabrications: none.**** All 10 cited key_papers plus both inline arXiv references (2106.14118; and the relevant 2605.10149) resolve to real papers with correct IDs, authors, and venues. The two highest-risk items ? the auth-gated Springer DOIs (10.1007/978-3-031-69073-0_15 and

10.1007/s42979-025-03935-0) ? are real, and their specific quantitative claims (e.g., BiLSTM service ~95%, end-of-rally via 780% consecutive no-motion frames over 187 rallies) verify exactly. Only cosmetic imperfections found: the SoccerNet paper's title is truncated (real title adds ": Submission to the SoccerNet Challenge 2022"), and BMN's author "Wen Sun" is shortened to "Wen."

****Verdict: held at partial-overlap.**** The survey is already conservative and self-honest (it calls its scarce-label badminton instantiation "novel-combination, not a clean gap"). My refutation searches confirmed that framing but surfaced closer prior art it missed, which nudges two sub-claims further toward "reinventing":

1. ****Reconciler**** ? `arXiv:2605.10149` (Improving TAS via Constraint-Aware Decoding, ICPR 2026) is a tighter analog than Activity Grammars: data-derived structural priors (transition confidence, boundary sets, per-class duration) resolved by global Viterbi as post-processing on a frozen model ? exactly the survey's own recommended upgrade, already published.
2. ****Overall assembly**** ? racket-sports rally/play-break segmentation via multi-cue fusion + lightweight classifier + rule/smoothing post-correction is ~20 years old (Kijak tennis play/break 2003; SmartTennisTV arXiv:1801.01430; HOG+ ℓ^2 -SVM+smoothing). The honest contribution narrows to the physics-cue taxonomy + interpretable logistic scorer + report-first/A-B-gated evidence-vs-decision reconciler + LOVO-by-match governance, not any single mechanism.

Course corrections (adopt 2605.10149 and the racket-sports lineage as citations; don't over-claim novelty over the badminton verdict-coders; close the search-to-citation coverage gap) are in the structured output.