

claude-windows / 7d6bbb83-42c0-459b-9e23-e14e96d4cb2d

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\subagents\workflows\wf_1dc2acd6-052\agent-af0d1649554d6f3c2.jsonl`  
- SHA-256: `ce23add0ab9f868f951ac2f8b171b9f68193003beb6f8b8d31274b4923581fed`  
- Source modified: `2026-06-13T20:24:08+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `wf_1dc2acd6-052`  
- session_id: `7d6bbb83-42c0-459b-9e23-e14e96d4cb2d`
```

Transcript

1. user (2026-06-13T20:19:41.354Z)

You are an ADVERSARIAL fact-checker. Theme: "Weakly/interval-supervised temporal localization + small-data evaluation protocols (LOVO, tIoU)".

Here are a teammate's survey findings (JSON):

```
{  
  "theme": "Weakly/interval-supervised temporal localization + small-data evaluation protocols (LOVO, tIoU)",  
  "our_pattern": "We segment a long static-camera badminton match into rally [start,end] intervals using scarce, expensive, interval-level golden labels (~6 matches). Multiple cheap \"producer\" cues each emit a (value, confidence) feature column over candidate windows; a tiny logistic regression learns to weight these columns and decides which windows are real rallies ? cues are never hand-coded if/else branches (\"signal not verdict\"). Expensive frame detection (shuttle CNN, person detector, audio hit detector) is persisted once and the cheap windowing+scorer re-scores deterministically offline (\"detection ? decision\"). A post-hoc \"reconciler\" linter flags where the decision contradicts persisted evidence (segment whose shuttle went inactive, zero net-crossings, over-merged segment with internal dead-time) and proposes A/B-measured corrections. New cues land flag-gated/default-OFF, graded vs golden labels with leave-one-VIDEO-out (group by match) using temporal IoU, promoted only on measured lift.",  
  "key_papers": [  
    {  
      "title": "Snorkel: Rapid Training Data Creation with Weak Supervision (Data Programming)",  
      "authors": "Ratner, Bach, Ehrenberg, Fries, Wu, Re",  
      "year": "2017",  
      "url_or_id": "arXiv:1711.10160",  
      "venue": "VLDB 2017",  
      "what_it_does": "Users write 'labeling functions' (arbitrary heuristics with unknown accuracy/correlation); a generative 'label model' denoises and combines them WITHOUT ground truth, learning a weight/accuracy per source from their agreements and disagreements, then trains an end discriminative model.",  
      "overlap": "strong",  
      "overlap_explanation": "This is the canonical formalization of our 'signal not verdict / no special-casing' core: each cue = a noisy weak-supervision source (their labeling function), and a learned model weights them instead of hand-coded branches. Key differences: (1) Snorkel learns source weights with NO labels via a generative model over LF agreement; we instead supervise our logistic regression directly with scarce golden intervals (closer to standard supervised feature stacking). (2) Snorkel operates at the data-point label level; our scorer operates at the candidate-window level for temporal localization. Their terminology (labeling function, label model, denoising, LF accuracy/correlation) is the established vocabulary we should adopt."},  
    {  
      "title": "SF-Net: Single-Frame Supervision for Temporal Action Localization",  
      "authors": "Ma, Yoon, Kim, Lee, Kang, Chen, Lim, Wang, Kira et al.",  
      "year": "2020",  
      "url_or_id": "arXiv:2003.06845",  
      "venue": "ECCV 2020",  
      "what_it_does": "Trains a temporal action localizer from a SINGLE timestamp per action instance; expands the labeled frame into pseudo-action segments via 'action frame expansion' and mines 'pseudo background frames' as negatives; evaluated with mAP@tIoU on THUMOS14/ActivityNet.",  
      "overlap": "partial",  
      "overlap_explanation": "Directly the literature for our regime: localizing temporal intervals from cheap, sparse annotations. Same metric family (mAP@tIoU). Difference: SF-Net learns a deep snippet model end-to-end from single frames and mines pseudo-labels; we hand-engineer scale-invariant window features and fit a tiny logistic regression on full interval labels. Confirms our problem is 'point/timestamp-supervised TAL', a named subfield ? we should use that terminology and its metric."},  
    {  
      "title": "Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses",  
      "authors": "Ju, Zhao, et al.",  
      "year": "2020",  
      "url_or_id": "arXiv:2012.08236",  
      "venue": "arXiv"
```

(2020)", "what_it_does": "Localizes actions from one timestamp per instance using a proposal-based paradigm that gives a more constrained solution space and consistent neighboring-frame predictions, vs frame-level paradigms whose sparse single-frame labels create a large solution space.", "overlap": "partial", "overlap_explanation": "Reinforces the 'point/timestamp supervision' framing for our scarce labels and argues proposal-based (candidate-window) prediction is better-constrained than dense frame prediction ? which is exactly our windowing-then-scoring design. We are doing a proposal/candidate-window approach without citing this lineage."}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries (ASRF)", "authors": "Ishikawa, Kasai, Aoki, Kataoka", "year": "2020", "url_or_id": "arXiv:2007.06866", "venue": "WACV 2021", "what_it_does": "Adds a Boundary Regression Branch that predicts action-boundary probability and uses detected boundaries to refine/relabel the per-frame Action Segmentation Branch output, cutting over-segmentation; reports segmental F1@k and edit distance (+13.7%/+16.1%).", "overlap": "partial", "overlap_explanation": "The closest method-level analog of our RECONCILERS: a separate signal (boundaries / persisted evidence) used post-hoc to correct an initial segmentation ? split over-merged segments and tighten boundaries. Differences: ASRF is a learned in-network branch and silently refines; our reconciler is a report-first, rule-precedence, A/B-gated, human-in-the-loop linter over persisted evidence. Their over-segmentation framing, boundary-refinement, and segmental F1@k / edit-distance metrics are directly relevant."}, {"title": "Structure analysis of soccer video with Hidden Markov Models (play/break segmentation)", "authors": "Xie, Xu, Chang, Divakaran, Sun", "year": "2001-2004", "url_or_id": "https://www.ee.columbia.edu/ln/dvmm/publications/01/icme01_soccer.pdf", "venue": "ICME 2001 / IEEE TMM", "what_it_does": "Defines two mutually-exclusive game states 'play' and 'break' from the rules of soccer; extracts domain features (dominant grass-color ratio, motion intensity) and uses HMM + dynamic programming to segment the match into play/break, ~86.5% segmentation accuracy.", "overlap": "strong", "overlap_explanation": "The seminal sports-structure analog of rally (play) vs dead-time (break) segmentation, and the origin of using domain cues + a temporal model rather than special-casing events. Key difference and lesson: it uses hand-tuned features in an HMM/generative pipeline (NOT a learned scorer over interchangeable feature columns) and per-frame state classification, where we use candidate windows + logistic regression. Establishes 'play/break segmentation' as the field term for exactly our task in a sibling sport."}, {"title": "SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos (+ 2023-2025 Challenges)", "authors": "Giancola, Amine, Dghaily, Ghanem; Cioppa et al. (challenge reports)", "year": "2018-2025", "url_or_id": "arXiv:1804.04527 (dataset); arXiv:2508.19182 (2025 results)", "venue": "CVPR Workshops / arXiv", "what_it_does": "Defines the action-spotting task and its evaluation: Average-Precision within a temporal tolerance delta (mAP@delta), with 'tight' (delta 1-5s) and 'loose' (5-60s) regimes; the de-facto sports-video temporal-localization benchmark and metric standard.", "overlap": "weak", "overlap_explanation": "Not the same task (point spotting vs interval segmentation) and a much larger labeled corpus, so weak methodological overlap. But it is THE established sports-temporal-localization evaluation standard; its tolerance-based mAP@delta and tight/loose reporting are conventions reviewers will expect alongside our tIoU, and worth reporting for boundary-slack sensitivity."}, {"title": "Distributional bias compromises leave-one-out cross-validation", "authors": "(2024/2025 paper, PMC12662204)", "year": "2024", "url_or_id": "https://pmc.ncbi.nlm.nih.gov/articles/PMC12662204/", "venue": "PMC (Briefings in Bioinformatics / GigaScience-class)", "what_it_does": "Shows LOO-CV induces a negative correlation between a fold's training-label mean and its held-out label; models regress to the training mean, so LOOCV systematically biases performance estimates and hyperparameter selection (a logistic regression scored ~0.23 auROC on random data). Proposes Rebalanced LOOCV (RLOOCV): drop an extra opposite-label training sample per fold to hold class balance constant.", "overlap": "strong", "overlap_explanation": "Directly targets our exact stack: leave-one-out cross-validation + logistic regression + tiny N. The bias is worst for small datasets and strong regularization ? precisely our 6-match, regularized-logistic regime. This is a concrete pitfall and a concrete fix (RLOOCV / class-balance-preserving folds) we should apply to our LOVO protocol and any threshold/lift comparisons."}, {"title": "Weakly Supervised TAL Through Learning Explicit Subspaces for Action and Context (action-context confusion)", "authors": "Liu, Wang, Tang, et al.", "year": "2021", "url_or_id": "arXiv:2103.16155", "venue": "AAAI 2021", "what_it_does": "Names and tackles 'action-context confusion': video-level supervision makes models latch onto co-occurring static context (e.g., the baseball field) instead of the action, hurting boundaries/completeness; learns separate action vs context subspaces.", "overlap": "weak", "overlap_explanation": "Methodologically distant (deep subspace learning) but the documented FAILURE MODE maps cleanly onto our setting: the static badminton court / warm-up context co-occurs with rallies, so a scorer can learn 'court is visible' rather than 'a point is in play.' This justifies our scale/translation-invariant, court-agnostic features and our warm-up/session-scope analyzer, and is the canonical name for the trap to

cite.]], "established_terms": ["Point / single-timestamp / point-level supervised Temporal Action Localization (the named subfield for our scarce interval/point labels)", "Weakly-supervised Temporal Action Localization (WTAL/WS-TAL) ? video-level labels", "Temporal Action Segmentation (the frame-labeling task; source of the over-segmentation literature)", "Play/Break segmentation (the sports-structure term for rally vs dead-time; from soccer-video analysis)", "Action spotting + tolerance-based mAP@delta (SoccerNet sports-temporal-localization metric standard)", "Data programming / labeling functions / label model / denoising (Snorkel) ? the established vocabulary for 'each cue is a noisy signal, a learned model weights them'", "Action-context confusion / completeness modeling (the WTAL failure-mode vocabulary)", "Leave-one-subject-out (LOSO) cross-validation ? the established name; 'leave-one-video-out (LOVO)' is our group-by-match instance", "mAP@tIoU and segmental F1@k / segmental edit distance (the metrics the field reports for interval localization and segmentation)"], "techniques_to_adopt": [{"technique": "Adopt Snorkel/data-programming vocabulary and consider a no-label 'label model' as a sanity check: frame each cue as a labeling function with an estimated accuracy/correlation; report per-cue accuracy and inter-cue correlation, and optionally fit a generative label model over cue agreements as a baseline against the supervised logistic regression.", "why": "Gives us the established terminology (labeling function, label model, denoising, source accuracy/correlation) for our 'signal not verdict' design, and the correlation diagnostic catches redundant/correlated cues ? a known weak-supervision failure that naive logistic stacking ignores.", "effort": "medium", "from_paper": "Snorkel: Rapid Training Data Creation with Weak Supervision (arXiv:1711.10160)"}, {"technique": "Switch LOVO scoring/thresholding to Rebalanced LOOCV (RLOOCV) or otherwise hold class balance constant across folds; never compare lift on raw LOO folds, and treat any near-random auROC swings as possible distributional-bias artifacts.", "why": "With only ~6 match-groups and a regularized logistic regression, standard LOO induces a documented inverse train-mean/test-label correlation that biases both the performance estimate and threshold/hyperparameter selection ? the exact knobs our experiment harness tunes when promoting a cue.", "effort": "low", "from_paper": "Distributional bias compromises leave-one-out cross-validation (PMC12662204)"}, {"technique": "Report tolerance-based metrics alongside tIoU: add SoccerNet-style mAP@delta in a tight (1-5s) and loose (5-60s) regime, and explicitly report boundary-slack sensitivity (how rally-start/-end precision changes the score).", "why": "tIoU alone hides where errors are (boundary slack vs whole-segment misses); the tolerance sweep is the convention reviewers expect for sports temporal localization and directly measures the serve-contact boundary-tightening our reconciler does.", "effort": "low", "from_paper": "SoccerNet: A Scalable Dataset for Action Spotting (arXiv:1804.04527; 2025 results arXiv:2508.19182)"}, {"technique": "Frame the reconciler's split/merge/boundary-tighten rules against the ASRF boundary-refinement formulation and report segmental F1@k + edit distance, not only tIoU; treat 'zero net-crossing / shuttle-inactive' as boundary/over-segmentation evidence analogous to a boundary-regression signal.", "why": "ASRF shows boundary detection is the standard remedy for over-segmentation and that segmental F1@k / edit distance are the metrics that expose merge/split errors which IoU can mask ? exactly the errors our reconciler targets.", "effort": "medium", "from_paper": "Alleviating Over-segmentation Errors by Detecting Action Boundaries / ASRF (arXiv:2007.06866)"}, {"technique": "Adopt a proposal/candidate-window paradigm justification and the 'play/break' framing explicitly; cite point-level TAL's argument that proposal-based prediction constrains the solution space better than dense frame prediction under sparse labels.", "why": "Validates our windowing-then-scoring choice with prior art and connects our task to the established play/break-segmentation lineage, improving comparability and reviewer legibility.", "effort": "low", "from_paper": "Point-Level Temporal Action Localization (arXiv:2012.08236) + Soccer HMM play/break (Xie/Xu, ICME 2001)"}, {"technique": "Make court/context-invariance an explicit, tested design goal: keep features scale/translation-invariant and court-agnostic, and add a probe that the scorer is not just keying on 'court visible' (e.g., evaluate on warm-up/practice stretches that share the court but lack point cadence).", "why": "Action-context confusion is the documented trap: with co-occurring static context the model learns the background, not the action. Our warm-up session-scope analyzer is the right instinct; the literature gives it a name and a way to test for the failure.", "effort": "medium", "from_paper": "WTAL Through Learning Explicit Subspaces for Action and Context (arXiv:2103.16155)"}], "pitfalls": ["L00/LOVO distributional bias: with ~6 folds and a regularized logistic regression, standard leave-one-out induces an inverse correlation between each fold's training-label mean and its held-out label, so the model regresses to the train mean and the CV score (and any threshold/lift it drives) is systematically biased ? can even fall below chance on signal-poor folds. Mitigate with class-balance-preserving folds (RLOOCV).", "Tiny-N variance / unstable promotion gate: 6 match-groups means each LOVO fold's metric has huge variance; a cue can show 'lift' from one lucky match. Without confidence intervals / significance testing across folds, the no-regression promotion rule will promote noise. Report per-fold spread, not just the mean.", "Action-context confusion: the static court and warm-up footage co-occur with rallies, so the scorer can learn 'court/players visible'

instead of 'point in play,' inflating in-distribution scores and failing on practice/dead-time that shares the background. (Documented WTAL failure mode.)", "Correlated cues break naive feature stacking: logistic regression assumes the contribution of each column is separable; highly correlated producers (e.g., shuttle-in-air and net-crossing) get double-counted or unstable weights. Weak-supervision literature explicitly models LF correlation; we currently don't.", "Over-segmentation / over-merging that tIoU hides: a single predicted segment merging two rallies across dead-time, or one rally split in two, can still score moderate tIoU. Segmental F1@k / edit-distance and a tolerance sweep are needed to expose these ? the precise errors the reconciler exists to fix.", "Reconciler that silently rewrites = label/eval leakage and metric gaming: post-hoc corrections tuned on (or evaluated against) the same scarce golden set can overfit the 6 matches. Our report-first + A/B-lift + LOVO gating is the right guard, but the lift itself must be measured under the bias-corrected CV above, or the 'measured lift' is illusory.", "Boundary/threshold tuning on the test fold: choosing tIoU thresholds, confidence cutoffs, or slack on data that includes the held-out video leaks. Group-by-match LOVO must wrap ALL tunable knobs (windowing, NMS/merge thresholds, scorer regularization), not just the final scorer fit."], "verdict": "partial-overlap", "verdict_rationale": "No single paper is our system, but every component of this theme is well-established and we should cite/reuse rather than claim novelty on the primitives. (1) 'Signal not verdict / learned scorer over cue columns' is data programming / weak supervision (Snorkel) ? the established formalization, with the caveat that we supervise directly on scarce labels rather than denoising label-free, and we ignore their inter-source-correlation modeling. (2) 'Scarce interval labels + temporal localization' is point/timestamp-supervised TAL (SF-Net, point-level TAL) evaluated with mAP@tIoU ? a named subfield. (3) 'Rally vs dead-time segmentation' is classic play/break structure analysis (Xie/Xu soccer HMM) ? a direct sibling-sport analog, and a recent badminton 'Point Segmentation' line exists. (4) The reconciler's over-segmentation/boundary correction mirrors ASRF-style boundary refinement, with segmental F1@k the expected metric. Where we are genuinely differentiated (a novel COMBINATION, not a gap): the 'detection ? decision / persist once, re-score forever' separation, the REPORT-FIRST, A/B-gated, idempotent, human-in-the-loop reconciler-as-linter (vs the field's silent in-network refinement), and the flag-gated default-OFF no-regression experiment harness on a deliberately tiny golden set. The most important course corrections are defensive on EVALUATION, not architecture: our LOVO + tiny-N + regularized-logistic stack sits squarely in the documented distributional-bias failure zone, so adopting RL00CV/balanced folds, per-fold variance reporting, tolerance-based metrics, and an action-context-confusion probe matters more than any architectural change.", "search_queries_used": ["weakly-supervised temporal action localization video-level labels survey", "point supervision temporal action localization single timestamp", "leave-one-subject-out cross-validation small-N sports machine learning overfitting", "temporal IoU evaluation action detection metric mAP tIoU", "badminton rally detection video segmentation broadcast play break shuttlecock", "tennis match play-break detection rally segmentation audio visual cues", "action segmentation over-segmentation refinement post-processing smoothing temporal", "multimodal late fusion calibration features logistic regression action recognition small data", "pseudo-label refinement consistency between prediction and evidence weak supervision verification", "sports highlight detection structure analysis play break domain knowledge features classifier soccer", "SoccerNet action spotting evaluation temporal tolerance metric ball-action", "data programming weak supervision label model combine noisy labeling functions Snorkel", "weakly supervised action localization action-context confusion completeness pitfall background"]}]

Your job is to be skeptical and protect against two failure modes: (1) FABRICATED citations, (2) OVERCLAIMED novelty.

DO, using WebSearch/WebFetch (load via ToolSearch "select:WebSearch,WebFetch" if needed):

- For each cited paper, search its EXACT title (+authors) to confirm it is a REAL paper. Mark `verified_real`
yes/no/uncertain and flag any that don't resolve in `fabrication_flags`. Be strict: a plausible-sounding title that you cannot find is a fabrication flag.
- For every "novel-combination" or "genuine-gap" verdict, actively try to REFUTE it: run fresh searches for the CLOSEST prior art that would downgrade the verdict toward "reinventing"/"partial-overlap". Default to assuming prior art exists until your searches come up empty. List any `closer_prior_art` you find (with real ids).
- Give a `corrected_verdict` + rationale. Add `extra_course_corrections` if your searches surfaced adopt/avoid notes.

Return the structured object.

2. assistant (2026-06-13T20:19:44.151Z)

I'll fact-check this survey adversarially. Let me start by loading the web tools and verifying each cited paper.

3. assistant (2026-06-13T20:19:48.278Z)

I'll verify each of the 9 cited papers and then adversarially probe the novelty claims. Let me batch the initial title-verification searches.

4. user (2026-06-13T20:19:57.773Z)

Web search results for query: "SF-Net Single-Frame Supervision for Temporal Action Localization ECCV 2020"

Links: [{"title": "SF-Net: Single-Frame Supervision for Temporal Action Localization | Springer Nature Link", "url": "https://link.springer.com/chapter/10.1007/978-3-030-58548-8_25"}, {"title": "SF-Net: Single-Frame Supervision for Temporal Action Localization · Issue #1877 · arXivTimes/arXivTimes", "url": "https://github.com/arXivTimes/arXivTimes/issues/1877"}, {"title": "[2003.06845v3] SF-Net: Single-Frame Supervision for Temporal Action Localization", "url": "https://arxiv.org/abs/2003.06845v3"}, {"title": "SF-Net: Single-Frame Supervision for Temporal Action Localization", "url": "https://www.ecva.net/papers/eccv_2020/papers_ECCV/papers/123490409.pdf"}, {"title": "SF-Net: Single-Frame Supervision for Temporal Action Localization | DeepAI", "url": "https://deepai.org/publication/sf-net-single-frame-supervision-for-temporal-action-localization"}, {"title": "Rethinking Pseudo-Label Guided Learning for Weakly Supervised Temporal Action Localization from the Perspective of Noise Correction", "url": "https://arxiv.org/pdf/2501.11124"}, {"title": "[2003.06845] SF-Net: Single-Frame Supervision for Temporal Action Localization", "url": "https://arxiv.org/abs/2003.06845"}, {"title": "Action-Agnostic Point-Level Supervision for Temporal Action Detection", "url": "https://arxiv.org/pdf/2412.21205"}, {"title": "Temporal Action Segmentation from Timestamp Supervision", "url": "https://arxiv.org/pdf/2103.06669"}, {"title": "RefineLoc: Iterative Refinement for Weakly-Supervised Action Localization", "url": "https://arxiv.org/pdf/1904.00227"}]

Based on the search results, here's information about the SF-Net paper:

Overview

SF-Net studies an intermediate form of supervision called single-frame supervision for temporal action localization (TAL), where annotators identify only a single frame within the temporal window of an action. This significantly reduces the labor cost compared to full supervision that requires annotating action boundaries, while introducing extra temporal action signals compared to weak supervision that only annotates video-level labels.

Key Technical Contributions

The approach proposes to predict an actionness score for each video frame, which along with a typical category score provides information about potential action occurrence and aids temporal boundary refinement during inference.

SF-Net mines pseudo action and background frames by identifying pseudo action frames through adaptively expanding each annotated single frame to nearby contextual frames and mining pseudo background frames from unannotated frames across multiple videos.

Results

In experiments on THUMOS14, GTEA, and BE01D datasets, SF-Net significantly improves upon state-of-the-art weakly-supervised methods and achieves comparable results to its fully-supervised counterpart which requires more resource intensive annotations.

The code is available at <https://github.com/Flowerfan/SF-Net>.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

5. user (2026-06-13T20:19:59.345Z)

Web search results for query: "Snorkel Rapid Training Data Creation with Weak Supervision Ratner Bach VLDB 2017"

Links: [{"title": "Snorkel: Rapid Training Data Creation with Weak Supervision Alexander Ratner", "url": "https://cs.brown.edu/people/sbach/files/ratner-vldb17.pdf"}, {"title": "Snorkel: rapid training data creation with weak supervision: Proceedings of the VLDB Endowment: Vol 11, No 3", "url": "https://dl.acm.org/doi/10.14778/3157794.3157797"}, {"title": "Snorkel: rapid training data creation with weak supervision - PubMed", "url": "https://pubmed.ncbi.nlm.nih.gov/32214778/"}, {"title": "[PDF] Snorkel: Rapid Training Data Creation with Weak Supervision | Semantic Scholar", "url": "https://www.semanticscholar.org/paper/Snorkel:-Rapid-Training-Data-Creation-with-Weak-Ratner-Bach/18bc1d4271abe8dd6e16179cdb06524a4f396e16"}, {"title": "Snorkel: Rapid Training Data Creation with Weak Supervision | Stanford DAWN", "url": "https://dawnd9.sites.stanford.edu/publications/snorkel/snorkel-rapid-training-data-creation-weak-supervision"}, {"title": "An Extensible Multimodal Multi-task Object Dataset with Materials", "url": "https://arxiv.org/pdf/2305.14352"}, {"title": "Learning Hyper Label Model for Programmatic Weak Supervision", "url": "https://arxiv.org/pdf/2207.13545"}, {"title": "[1711.10160] Snorkel: Rapid Training Data Creation with Weak Supervision", "url": "https://arxiv.org/abs/1711.10160"}, {"title": "Statistically Valid Post-Deployment Monitoring Should Be Standard for AI-Based Digital Health", "url": "https://arxiv.org/pdf/2506.05701"}, {"title": "Label Propagation with Weak Supervision", "url": "https://arxiv.org/pdf/2210.03594"}]

I found the paper you're looking for! Here are the search results for "Snorkel: Rapid Training Data Creation with Weak Supervision" by Ratner, Bach, and colleagues:

Key Information

Labeling training data is increasingly the largest bottleneck in deploying machine learning systems. Snorkel is a first-of-its-kind system that enables users to train state-of-the-art models without hand labeling any training data.

Users write labeling functions that express arbitrary heuristics, which can have unknown accuracies and correlations. Snorkel denoises their outputs without access to ground truth by incorporating the first end-to-end implementation of the data programming machine learning paradigm.

Publication Details

The paper is authored by A. Ratner, S. H. Bach, H. Ehrenberg, J. Fries, S. Wu, and C. Ré, published in PVLDB, Volume 11, Issue 3 in 2017.

Available Resources

The paper is available from multiple sources:

- The official PDF at Brown University's CS department
- ACM Digital Library (DOI: 10.14778/3157794.3157797)
- arXiv preprint (arxiv:1711.10160)
- Stanford DAWN project website

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

6. assistant (2026-06-13T20:21:18.605Z)

Both verified. Snorkel (arXiv:1711.10160, VLDB 2017) confirmed real. SF-Net (arXiv:2003.06845, ECCV 2020) confirmed real. Note: the survey's "evaluated with mAP@tIoU on THUMOS14/ActivityNet" ? actual datasets are THUMOS14/GTEA/BE0ID, not ActivityNet. Minor inaccuracy, not a fabrication. Let me continue.

7. user (2026-06-13T20:21:23.970Z)

Web search results for query: ""Point-Level Temporal Action Localization" "Bridging Fully-supervised Proposals" weakly-supervised losses arXiv"

Links: [{"title": "[2012.08236v1] Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses", "url": "https://arxiv.org/abs/2012.08236v1"}, {"title": "[2012.08236] Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses", "url": "https://arxiv.org/abs/2012.08236"}, {"title": "Distilling Vision-Language Pre-training to Collaborate with Weakly-Supervised Temporal Action Localization", "url": "https://arxiv.org/pdf/2212.09335"}, {"title": "Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses - ADS", "url": "https://ui.adsabs.harvard.edu/abs/2020arXiv201208236J/abstract"}, {"title": "Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses | Request PDF", "url": "https://www.researchgate.net/publication/347300936_Point-Level_Temporal_Action_Localization_Bridging_Fully-supervised_Proposals_to_Weakly-supervised_Losses"}, {"title": "Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses", "url": "https://arxiv.org/pdf/2012.08236"}, {"title": "[PDF] Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses | Semantic Scholar", "url": "https://www.semanticscholar.org/paper/Point-Level-Temporal-Action-Localization:-Bridging-Ju-Zhao/fb932412e570f507f783e2ef633ea64677bd268e"}, {"title": "GitHub - zhenyingfang/Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation: Temporal Action Detection & Weakly Supervised Temporal Action Detection & Temporal Action Proposal Generation · GitHub", "url": "https://github.com/zhenyingfang/Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation"}]

I found the paper you're looking for. Here are the search results:

The paper "Point-Level Temporal Action Localization: Bridging Fully-supervised Proposals to Weakly-supervised Losses" addresses point-level temporal action localization (PTAL), which aims to localize actions in untrimmed videos with only one timestamp annotation for each action instance.

The paper explores the proposal-based prediction paradigm for point-level annotations, which has the advantage of more constrained solution space and consistent predictions among neighboring frames. The approach uses point-level annotations as keypoint supervision to train a keypoint detector, and introduces a mapper module that bridges the fully-supervised framework with weak supervision. The authors note this is the first work to leverage the fully-supervised paradigm for the point-level setting.

The paper was published on arXiv in December 2020 (arXiv:2012.08236) and authored by Chen Ju, Peisen Zhao, Ya Zhang, Yanfeng Wang, and Qi Tian from Shanghai Jiao Tong University and Huawei Cloud & AI.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

8. user (2026-06-13T20:21:27.497Z)

Web search results for query: "Alleviating Over-segmentation Errors by Detecting Action Boundaries ASRF Ishikawa Kasai Aoki Kataoka WACV"

Links: [{"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries | IEEE Conference Publication | IEEE Xplore", "url": "https://ieeexplore.ieee.org/document/9423184/"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries", "url": "https://yiskw713.github.io/asrf/"}], [{"title": "Alleviating Over-segmentation"}]

Errors by Detecting Action Boundaries", "url": "https://openaccess.thecvf.com/content/WACV2021/papers/Ishikawa_Alleviating_Over-Segmentation_Errors_by_Detecting_Action_Boundaries_WACV_2021_paper.pdf"}, {"title": "asrf/evaluate.py at main · yiskw713/asrf", "url": "https://github.com/yiskw713/asrf/blob/main/evaluate.py"}, {"title": "Alleviating Over-Segmentation Errors by Detecting Action Boundaries", "url": "https://papertalk.org/papertalks/27315"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries", "url": "https://arxiv.org/pdf/2007.06866"}, {"title": "GitHub - yiskw713/asrf: Code for 'Alleviating Over-segmentation Errors by Detecting Action Boundaries' accepted in WACV2021", "url": "https://github.com/yiskw713/asrf"}, {"title": "Alleviating Over-segmentation Errors by Detecting Action Boundaries | Request PDF", "url": "https://www.researchgate.net/publication/342944719_Alleviating_Over-segmentation_Errors_by_Detecting_Action_Boundaries"}, {"title": "WACV 2021 Open Access Repository", "url": "https://openaccess.thecvf.com/content/WACV2021/html/Ishikawa_Alleviating_Over-Segmentation_Errors_by_Detecting_Action_Boundaries_WACV_2021_paper.html"}]

I found the paper! Here's what I found about this work:

Paper Information

The paper "Alleviating Over-segmentation Errors by Detecting Action Boundaries" was authored by Yuchi Ishikawa, Seito Kasai, Yoshimitsu Aoki, and Hirokatsu Kataoka, and published in WACV 2021.

Overview

The paper proposes an effective framework for the temporal action segmentation task, namely an Action Segment Refinement Framework (ASRF). The model architecture consists of a long-term feature extractor and two branches: the Action Segmentation Branch (ASB) and the Boundary Regression Branch (BRB). The long-term feature extractor provides shared features for the two branches with a wide temporal receptive field. The ASB classifies video frames with action classes, while the BRB regresses the action boundary probabilities.

Key Contributions

The framework divides temporal action segmentation into frame-wise action classification and action boundary regression, and refines frame-level hypotheses of action classes using predicted action boundaries. The model outperforms state-of-the-art methods on three challenging datasets, offering an improvement of up to 13.7% in terms of segmental edit distance and up to 16.1% in terms of segmental F1 score.

You can find the official code repository on [GitHub](https://github.com/yiskw713/asrf) and the paper on [arXiv](https://arxiv.org/pdf/2007.06866) and the [WACV 2021 proceedings](https://openaccess.thecvf.com/content/WACV2021/papers/Ishikawa_Alleviating_Over-Segmentation_Errors_by_Detecting_Action_Boundaries_WACV_2021_paper.pdf).

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

9. assistant (2026-06-13T20:21:31.741Z)

Both verified. Point-Level TAL (arXiv:2012.08236, authors Ju, Zhao, Zhang, Wang, Tian) confirmed ? survey lists "Ju, Zhao, et al." which is correct. ASRF (arXiv:2007.06866, WACV 2021, all four authors match, +13.7%/+16.1% figures match) confirmed real. Let me verify the soccer HMM, SoccerNet, the LOOCV bias paper, and the action-context confusion paper.

10. user (2026-06-13T20:21:41.860Z)

Web search results for query: "Xie Xu Chang Divakaran Sun structure analysis soccer video Hidden Markov Models play break segmentation ICME 2001"

Links: [{"title": "(PDF) Structure analysis of soccer video with hidden Markov models", "url": "https://www.academia.edu/8460278/Structure_analysis_of_soccer_video_with_hidden_Markov_models"}, {"title": "Unsupervised Mining of Statistical Temporal Structures in Video | Springer Nature Link", "url": "https://link.springer.com/chap/10.1007/978-3-319-25000-0_10"}]


```
url":"https://link.springer.com/chapter/10.1007/978-1-4757-6928-9_10"}, {"title": "Highlights modeling and detection in sports videos - Pattern Analysis and Applications", "url": "https://link.springer.com/article/10.1007/s10044-004-0234-1"}, {"title": "Multimodal Semantic Analysis and Annotation for Basketball Video | Journal on Advances in Signal Processing | Springer Nature Link", "url": "https://link.springer.com/article/10.1155/ASP/2006/32135"}, {"title": "Audiovisual integration for tennis broadcast structuring | Multimedia Tools and Applications | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s11042-006-0031-5"}, {"title": "Common Visual Cues for Sports Highlights Modeling | Multimedia Tools and Applications | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s11042-005-2575-1"}, {"title": "Automatic composition of broadcast sports video | Multimedia Systems | Springer Nature Link", "url": "https://link.springer.com/article/10.1007/s00530-008-0112-6"}, {"title": "US6865226B2 - Structural analysis of videos with hidden markov models and dynamic programming - Google Patents", "url": "https://patents.google.com/patent/US6865226?q=hidden+markov+model"}, {"title": "(PDF) Algorithms And System For Segmentation And Structure Analysis In Soccer Video", "url": "https://www.researchgate.net/publication/228907792_Algorithms_And_System_For_Segmentation_And_Structure_Analysis_In_Soccer_Video"}]
```

I found relevant information about your search query. Here are the results:

Xu, P., Xie, L., Chang, S.-F., Divakaran, A., Vetro, A., and Sun, H. presented "Algorithms and systems for segmentation and structure analysis in soccer video" at the IEEE International Conference on Multimedia and Expo (ICME) in Tokyo, Japan.

Additionally, Xie, Xu, Chang, Divakaran, and Sun also published work on "Structure analysis of soccer video with domain knowledge and hidden Markov models", which was presented at a later conference.

The work on heuristic rules based on domain specific features and dominant color ratios was used to segment play and break in soccer video, as described in the Xu et al. ICME 2001 paper.

These papers appear to be seminal works in the field of sports video analysis, specifically addressing segmentation and structure analysis using Hidden Markov Models and domain-specific features for detecting play/break sequences in soccer video.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

11. user (2026-06-13T20:21:41.897Z)

Web search results for query: "SoccerNet Scalable Dataset Action Spotting Soccer Videos Giancola Ghanem arXiv 1804.04527"

```
Links: [{"title": "[1804.04527v2] SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos", "url": "https://arxiv.org/abs/1804.04527v2"}, {"title": "SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos", "url": "https://arxiv.org/pdf/1804.04527"}, {"title": "[1804.04527] SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos", "url": "https://arxiv.org/abs/1804.04527"}, {"title": "[PDF] SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos | Semantic Scholar", "url": "https://www.semanticscholar.org/paper/SoccerNet:-A-Scalable-Dataset-for-Action-Spotting-Giancola-Amine/982f2025925062aeafac07ae015c9ed273e4d3d6"}, {"title": "(PDF) SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos", "url": "https://www.researchgate.net/publication/324492785_SoccerNet_A_Scalable_Dataset_for_Action_Spotting_in_Soccer_Videos"}, {"title": "SoccerNet: A Scalable Dataset for Action Spotting in Soccer Videos", "url": "https://openaccess.thecvf.com/content_cvpr_2018_workshops/papers/w34/Giancola_SoccerNet_A_Scalable_CVPR_2018_paper.pdf"}]
```

Based on the search results, here's information about the paper you're looking for:

SoccerNet is a benchmark for action spotting in soccer videos, composed of 500 complete soccer games from six main European leagues, covering three seasons from 2014 to 2017 and a total duration of 764 hours.

The dataset includes 6,637 temporal annotations automatically parsed from online match reports at a one minute resolution for three main classes of events (Goal, Yellow/Red Card, and Substitution), which are then manually refined to a one second resolution by anchoring them at a single timestamp following well-defined soccer rules.

With an average of one event every 6.9 minutes, this dataset focuses on the problem of localizing very sparse events within long videos. The authors' best model for classifying temporal segments of length one minute reaches a mean Average Precision (mAP) of 67.8%, while for the spotting task, their baseline reaches an Average-mAP of 49.7% for tolerances ranging from 5 to 60 seconds.

The paper was authored by Silvio Giancola, Mohieddine Amine, Tarek Dghaily, and Bernard Ghanem from King Abdullah University of Science and Technology (KAUST), and was published on arXiv on April 12, 2018 (arXiv:1804.04527).

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

12. assistant (2026-06-13T20:21:48.079Z)

Both verified. SoccerNet (arXiv:1804.04527, CVPR Workshops 2018, authors match) confirmed real. The soccer HMM paper: the search confirms a real lineage of Xie/Xu/Chang/Divakaran/Sun soccer-video play/break work with HMM + dominant grass-color + dynamic programming. The exact title "Structure analysis of soccer video with Hidden Markov Models" exists (academia.edu link), and there's a closely-related "Structure analysis of soccer video with domain knowledge and hidden Markov models." Real, with a slightly fuzzy title/year, which is acceptable given this is a 2001-2004 body of work. Let me verify the two most important and most specific claims: the LOOCV bias paper (PMC ID) and the action-context confusion AAAI paper, plus check the 2025 SoccerNet results ID.

13. user (2026-06-13T20:21:56.551Z)

Web search results for query: ""Distributional bias compromises leave-one-out cross-validation" Rebalanced LOOCV RL00CV"

Links: [{"title": "[Literature Review] Distributional bias compromises leave-one-out cross-validation", "url": "https://www.themoonlight.io/en/review/distributional-bias-compromises-leave-one-out-cross-validation"}, {"title": "Distributional bias compromises leave-one-out cross-validation | Science Advances", "url": "https://www.science.org/doi/10.1126/sciadv.adx6976"}, {"title": "Distributional bias compromises leave-one-out cross-validation", "url": "https://www.ncbi.nlm.nih.gov/pmc/articles/PMC12662204/"}, {"title": "Distributional bias compromises leave-one-out cross-validation", "url": "https://www.ovid.com/journals/sciad/fulltext/10.1126/sciadv.adx6976~distributional-bias-compromises-leave-one-out"}, {"title": "Recent paper: Distributional bias compromises leave-one-out cross-validation - Modeling - The Stan Forums", "url": "https://discourse.mc-stan.org/t/recent-paper-distributional-bias-compromises-leave-one-out-cross-validation/40743"}, {"title": "[2406.01652] Distributional bias compromises leave-one-out cross-validation", "url": "https://arxiv.org/abs/2406.01652"}]

Based on the search results, here's what I found about this topic:

Distributional Bias in Leave-One-Out Cross-Validation

Standard implementations of leave-P-out cross-validation (LPOCV) introduce distributional bias? a negative correlation, across folds, between the average label of the corresponding training and test sets. In LOOCV, this approach inadvertently introduces a negative correlation between the average label of the training set and the label of the test instance.

This phenomenon can be demonstrated via a simple predictor that takes advantage of this bias, and in practice, due to the tendency of machine learning models to regress to the mean of the training set labels, this bias tends to decrease common performance evaluation metrics, such as auROC, auPR, or R^2 .

Rebalanced LOOCV (RLOOCV)

To address this issue, researchers proposed Rebalanced LOOCV (RLOOCV) as a correction. In RLOOCV, for each individual instance that is left out during cross-validation, a sample that has an opposite label to the left-out instance is also omitted from the training set, maintaining a consistent label balance across training instances.

RLOOCV maintains i.i.d. distributions of training and test sets, while LOOCV does not, which ensures that RLOOCV cannot erroneously overinflate performance estimates. This correction consistently increases measured machine learning performances under multiple machine learning benchmarks.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

14. user (2026-06-13T20:22:00.501Z)

Web search results for query: "Weakly Supervised Temporal Action Localization Learning Explicit Subspaces Action Context confusion AAAI 2021 arXiv 2103.16155"

Links: [{"title": "[2103.16155] Weakly Supervised Temporal Action Localization Through Learning Explicit Subspaces for Action and Context", "url": "https://arxiv.org/abs/2103.16155"}, {"title": "Weakly Supervised Temporal Action Localization Through Learning Explicit Subspaces for Action and Context | Proceedings of the AAAI Conference on Artificial Intelligence", "url": "https://ojs.aaai.org/index.php/AAAI/article/view/16323"}, {"title": "Weakly Supervised Temporal Action Localization Through Learning Explicit Subspaces for Action and Context", "url": "https://arxiv.org/pdf/2103.16155"}, {"title": "Weakly Supervised Temporal Action Localization Through ...", "url": "https://cdn.aaai.org/ojs/16323/16323-13-19817-1-2-20210518.pdf"}, {"title": "GitHub - zhenyingfang/Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation: Temporal Action Detection & Weakly Supervised Temporal Action Detection & Temporal Action Proposal Generation · GitHub", "url": "https://github.com/zhenyingfang/Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation"}, {"title": "Background-Click Supervision for Temporal Action Localization", "url": "https://arxiv.org/pdf/2111.12449"}, {"title": "Fine-grained Temporal Contrastive Learning for Weakly-supervised Temporal Action Localization", "url": "https://arxiv.org/pdf/2203.16800"}, {"title": "Improving Weakly Supervised Temporal Action Localization by Bridging Train-Test Gap in Pseudo Labels", "url": "https://arxiv.org/pdf/2304.07978"}, {"title": "Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation/README.md at master · zhenyingfang/Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation", "url": "https://github.com/zhenyingfang/Awesome-Temporal-Action-Detection-Temporal-Action-Proposal-Generation/blob/master/README.md"}]

I've successfully found the paper you're looking for. Here's the information:

Paper Summary

The paper "Weakly Supervised Temporal Action Localization Through Learning Explicit Subspaces for Action and Context" was authored by Liu, Z., Wang, L., Tang, W., Yuan, J., Zheng, N., & Hua, G. and published in the Proceedings of the AAAI Conference on Artificial Intelligence, volume 35, pages 2242-2250 in 2021. The paper is also available on arXiv as arXiv:2103.16155.

Key Contributions

The paper addresses Weakly-supervised Temporal Action Localization (WS-TAL), which learns to localize temporal starts and ends of action instances in a video under only video-level supervision.

The main challenge the paper tackles is that existing WS-TAL methods cannot distinguish frequently co-occurring contextual background from actual action instances?a problem termed "action-context confusion".

To address this challenge, the paper introduces a framework that learns two feature subspaces

respectively for actions and their context, allowing action instances to be localized more precisely without distraction from the context.

The proposed approach outperforms state-of-the-art WS-TAL methods on three benchmarks: THUMOS14, ActivityNet v1.2 and v1.3 datasets.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

15. assistant (2026-06-13T20:22:06.601Z)

Both verified. The L00CV bias paper is REAL but the survey has a notable inaccuracy: it's published in **Science Advances** (DOI 10.1126/sciadv.adx6976, arXiv:2406.01652), NOT "Briefings in Bioinformatics / GigaScience-class" as the survey claims. The PMC ID PMC12662204 resolves correctly. The RL00CV fix and the negative-correlation mechanism are accurately described. The action-context confusion paper (arXiv:2103.16155, AAAI 2021, authors Liu/Wang/Tang/Yuan/Zheng/Hua) confirmed real.

Now let me verify the 2025 SoccerNet results ID and then probe the novelty claims adversarially.

16. user (2026-06-13T20:22:15.293Z)

Web search results for query: "arXiv 2508.19182 SoccerNet 2025 challenge results"

Links: [{"title": "SoccerNet 2025 Challenges Results - ADS", "url": "https://ui.adsabs.harvard.edu/abs/2025arXiv250819182G/abstract"}, {"title": "(PDF) SoccerNet 2025 Challenges Results", "url": "https://www.researchgate.net/publication/394978601_SoccerNet_2025_Challenges_Results"}, {"title": "SoccerNet 2025 Challenges Results | alphaXiv", "url": "https://www.alphaxiv.org/overview/2508.19182"}, {"title": "[2508.19182] SoccerNet 2025 Challenges Results", "url": "https://arxiv.org/abs/2508.19182"}, {"title": "[Literature Review] SoccerNet 2025 Challenges Results", "url": "https://www.themoonlight.io/en/review/soccernet-2025-challenges-results"}, {"title": "SoccerNet on X: \"? The SoccerNet Challenge 2025 Report is out! ? ? https://t.co/Dkjp53ibf Thanks to all teams pushing the limits of AI in football ?? Stay tuned: SoccerNet Challenges 2026 will be unveiled progressively in the next 2 months! Join us on Discord ? https://t.co/0WpRf2UThn\" / X", "url": "https://x.com/soccernet_org/status/1960676485410197954"}, {"title": "FOOTPASS: A Multi-Modal Multi-Agent Tactical Context Dataset for Play-by-Play Action Spotting in Soccer Broadcast Videos", "url": "https://arxiv.org/pdf/2511.16183"}]

I found information about the SoccerNet 2025 Challenges Results paper (arXiv 2508.19182). Here's what the search revealed:

The SoccerNet 2025 Challenges mark the fifth annual edition of the SoccerNet open benchmarking effort, dedicated to advancing computer vision research in football video understanding.

The challenges span four vision-based tasks: (1) Team Ball Action Spotting, focused on detecting ball-related actions in football broadcasts and assigning actions to teams; (2) Monocular Depth Estimation, targeting the recovery of scene geometry from single-camera broadcast clips through relative depth estimation for each pixel; (3) Multi-View Foul Recognition, requiring the analysis of multiple synchronized camera views to classify fouls and their severity; and (4) Game State Reconstruction, aimed at localizing and identifying all players from a broadcast video to reconstruct the game state on a 2D top-view of the field.

Across all tasks, participants were provided with large-scale annotated datasets, unified evaluation protocols, and strong baselines as starting points.

The challenges saw significant participation, with over 280 participants across all tasks, consistently outperforming the proposed baselines and introducing novel methodologies.

The paper was released on August 26, 2025, and presents the results of each challenge, highlights the top-performing solutions, and provides insights into the progress made by the community.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

17. user (2026-06-13T20:22:16.662Z)

Web search results for query: "badminton rally segmentation point segmentation broadcast video temporal weakly supervised"

Links: [{"title": "Towards Real-Time Analysis of Broadcast Badminton Videos", "url": "https://arxiv.org/pdf/2308.12199"}, {"title": "(PDF) Analyzing Racket Sports From Broadcast Videos", "url": "https://www.academia.edu/70243184/Analyzing_Racket_Sports_From_Broadcast_Videos"}, {"title": "TennisExpert: Towards Expert-Level Analytical Sports Video Understanding", "url": "https://arxiv.org/pdf/2603.13397"}, {"title": "BFMD: A Full-Match Badminton Dense Dataset for Dense Shot Captioning", "url": "https://arxiv.org/pdf/2603.25533"}, {"title": "(PDF) Towards Structured Analysis of Broadcast Badminton Videos", "url": "https://www.researchgate.net/publication/322075902_Towards_Structured_Analysis_of_Broadcast_Badminton_Videos"}, {"title": "Badminton Video Analysis based on Spatiotemporal and Stroke Features", "url": "https://www.researchgate.net/publication/317159692_Badminton_Video_Analysis_based_on_Spatiotemporal_and_Stroke_Features"}, {"title": "DOI: 10.1111/cgf.14786 COMPUTER GRAPHICS forum", "url": "https://diglib.eg.org/server/api/core/bitstreams/ad201b25-2386-4522-9ced-776e1656c745/content"}, {"title": "Exploration of Player Behaviors from Broadcast Badminton ...", "url": "https://people.cs.nyu.edu.tw/~yushuen/data/BadmintonVis23.pdf"}, {"title": "Bridging the Gap: Doubles Badminton Analysis with Singles-Trained Models", "url": "https://arxiv.org/html/2508.13507v1"}, {"title": "Towards Structured Analysis of Broadcast Badminton Videos", "url": "https://arxiv.org/pdf/1712.08714"}]

Based on the search results, I found several relevant research papers and resources on badminton rally and point segmentation in broadcast videos:

Key Research Findings

****Point/Rally Segmentation in Badminton Videos:****

Broadcast badminton matches are divided into gameplays and replays, where "rally" frames are frames where actual gameplay is happening and "non-rally" frames are replays or player rest periods. This classification is called "Point Segmentation," which allows researchers to keep all rally frames, group them into rallies according to points scored, and discard non-rally frames.

****Applications and Methods:****

Researchers propose methods to analyze badminton broadcast videos by segmenting the points played, tracking and recognizing the players in each point, and annotating their badminton strokes. State-of-the-art machine learning methods are applied to partition broadcast videos into segments where each segment shows a badminton rally.

****Recent Datasets:****

Recent full-match badminton datasets collect videos from official BWF World Tour tournaments and include frame-level temporal annotations segmenting broadcasts into rallies and broadcast interruption segments.

The research demonstrates an active area of study in automated video analysis for racket sports, with focus on temporal segmentation and understanding game structure from broadcast footage.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

18. assistant (2026-06-13T20:22:26.379Z)

The 2025 SoccerNet ID (arXiv:2508.19182) is real. Importantly, the badminton search confirms the survey's "Point Segmentation" claim ? "Towards Structured Analysis of Broadcast Badminton

Videos" (arXiv:1712.08714) defines exactly the rally/non-rally "Point Segmentation" task the survey cites. This is real prior art for the badminton-specific task and strengthens the "partial-overlap" verdict.

Now let me adversarially probe the genuine-differentiation claims. The survey claims novelty in the COMBINATION (not a gap), specifically: (1) "detection ? decision / persist once, re-score forever," (2) report-first/human-in-the-loop reconciler-as-linter, (3) flag-gated default-OFF no-regression harness on a tiny golden set. Let me search for prior art that would downgrade these.

19. user (2026-06-13T20:22:36.429Z)

Web search results for query: "decouple feature extraction from decision precomputed features re-score offline video action detection deterministic"

Links: [{"title": "Decoupling Localization and Classification in Single Shot Temporal Action Detection", "url": "https://arxiv.org/pdf/1904.07442"}, {"title": "Online Action Detection", "url": "https://arxiv.org/pdf/1604.06506"}, {"title": "Detection of human actions from video data", "url": "https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/9613276"}, {"title": "Colar: Effective and Efficient Online Action Detection by Consulting Exemplars", "url": "https://arxiv.org/pdf/2203.01057"}, {"title": "Linear-time Online Action Detection From 3D Skeletal Data Using Bags of Gesturelets", "url": "https://arxiv.org/pdf/1502.01228"}, {"title": "Leaving Some Stones Unturned: Dynamic Feature Prioritization for Activity Detection in Streaming Video | SpringerLink", "url": "https://link.springer.com/chapter/10.1007/978-3-319-46478-7_48"}, {"title": "Efficient Video Action Detection with Token Dropout and Context Refinement", "url": "https://openaccess.thecvf.com/content/ICCV2023/papers/Chen_Efficient_Video_Action_Detection_with_Token_Dropout_and_Context_Refinement_ICCV_2023_paper.pdf"}, {"title": "Visual Feature Learning on Video Object and Human Action Detection: A Systematic Review", "url": "https://www.mdpi.com/2072-666X/13/1/72"}, {"title": "Video-Based Recognition of Human Activity Using Novel Feature Extraction Techniques", "url": "https://www.mdpi.com/2076-3417/13/11/6856"}, {"title": "A Survey of Video Action Recognition Based on Deep Learning - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/pii/S0950705125006409"}]

Based on the web search results, I found several relevant papers and resources related to video action detection with decoupled feature extraction and offline processing approaches:

Key Findings:

****Decoupled Feature Extraction in Action Detection:****

Decoupled Single Shot temporal Action Detection (Decouple-SSAD) explores separable localization and classification mechanisms in one-stage action detection frameworks, studying the problem of learning specialized features useful for optimizing each subtask through two separate branches.

****Offline Video Action Detection:****

In offline detection, the goal is to find the start and end frame of any action that occurs in the video, with all information of the video available at once, and calculation time is not an issue.

****Pre-extracted Features Approach:****

Given a series of pre-extracted video features, systems can connect frames with historical neighbors and model local evolution while comparing frames with representative exemplars to model category particularity.

****Efficient Feature Re-scoring:****

Seq-NMS finds the most accurate region proposals from a video clip by selecting the highest score of a frame sequence from a video, rescores the sequence, and removes the overlapped region proposals.

****Keyframe-based Feature Extraction:****

Methods using keyframe selection and spatiotemporal feature extraction reduce computation while enhancing feature representations through decoupled convolution and feature fusion.

These approaches align with the concept of decoupling feature extraction from decision-making and using precomputed features for offline video action detection.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

20. user (2026-06-13T20:22:38.659Z)

Web search results for query: "human-in-the-loop verification rule-based linter inconsistency prediction evidence weak labels propose corrections A/B"

Links: [{"title": "Human-in-the-Loop, Human-on-the-Loop, and LLM-as-a-Judge for Validating AI Outputs", "url": "https://kili-technology.com/blog/human-in-the-loop-human-on-the-loop-and-llm-as-a-judge-for-validating-ai-outputs"}, {"title": "Detecting Corpus-Level Knowledge Inconsistencies in Wikipedia with Large Language Models", "url": "https://arxiv.org/pdf/2509.23233"}, {"title": "Human-in-the-Loop AI Use in Ongoing Process Verification in the Pharmaceutical Industry", "url": "https://www.mdpi.com/2078-2489/16/12/1082"}, {"title": "E2Edev: Benchmarking Large Language Models in End-to-End Software Development Task", "url": "https://arxiv.org/pdf/2510.14509"}, {"title": "Verification of design libraries and databases", "url": "https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/8949203"}, {"title": "Knowledge graph validation by integrating LLMs and human-in-the-loop - ScienceDirect", "url": "https://www.sciencedirect.com/science/article/pii/S030645732500086X"}, {"title": "Fine-Grained Detection of Context-Grounded Hallucinations Using LLMs", "url": "https://arxiv.org/pdf/2509.22582"}, {"title": "Towards Scientific Intelligence: A Survey of LLM-based Scientific Agents", "url": "https://arxiv.org/pdf/2503.24047"}]

I've performed the web search with your query. The results returned articles and papers related to human-in-the-loop systems, verification processes, and inconsistency detection. Here's what was found:

The search results include:

1. **Human-in-the-Loop verification approaches** - Resources describing human-in-the-loop assistants for inconsistency detection that involve both research (gathering relevant evidence) and verification (determining whether evidence contradicts a given fact).
2. **Knowledge graph validation** - Papers on state-of-the-art validation techniques categorized as automated methods relying on statistics or rules, and methods relying on human-in-the-loop such as crowdsourcing.
3. **LLM-based inconsistency detection** - Work on using LLMs to identify potential inconsistencies with high recall, followed by human annotators reviewing outputs to filter false positives.
4. **Expert oversight in agents** - Human-in-the-loop verification that integrates human domain experts into agent workflows to review outputs, provide approval/rejection decisions, and offer qualitative feedback.

However, the search results don't specifically contain papers matching all the specific keywords in your query (weak labels, propose corrections, A/B testing in the context of rule-based linters). If you're looking for a specific research paper or topic, you might want to refine the search query or provide additional context about what you're looking for.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

21. assistant (2026-06-13T20:22:44.810Z)

Useful. The "decouple feature extraction from decision" pattern has prior art (Decouple-SSAD, arXiv:1904.07442 ? decouples localization from classification; offline detection with pre-extracted features is standard). But the survey's specific framing ("persist expensive detection once, deterministically re-score with cheap windowing forever") is more of an engineering/MLOps

pattern than a claimed research contribution ? the survey itself classifies it as a novel *combination*, not a gap, which is defensible. The reconciler-as-linter search did not surface a direct analog in the video-segmentation space; the closest is ASRF (already cited) for the in-network silent version. Let me do two more targeted refutation searches: one on the experiment-harness / no-regression eval pattern, and one on pseudo-label/prediction-evidence consistency verification, to confirm whether the "novel combination" holds.

22. user (2026-06-13T20:22:54.783Z)

Web search results for query: "flag-gated feature shadow evaluation no-regression A/B testing machine learning model promotion offline guardrail"

Links: [{"title": "Minimize the production impact of ML model updates with Amazon SageMaker shadow testing | Artificial Intelligence", "url": "https://aws.amazon.com/blogs/machine-learning/minimize-the-production-impact-of-ml-model-updates-with-amazon-sagemaker-shadow-testing/"}, {"title": "Shadow Testing for AI: How to Evaluate Models Without Affecting Users", "url": "https://www.statsig.com/perspectives/shadow-testing-ai-model-evaluation"}, {"title": "Questions Answered ? Shadow Deployment In Machine Learning | by Demetrios Brinkmann | MLOps.community | Medium", "url": "https://medium.com/mlops-community/questions-answered-shadow-deployment-in-machine-learning-5ee5a8854e10"}, {"title": "GitHub - mohammedbilalkhan/sagemaker-deployment-guardrail-shadow-testing", "url": "https://github.com/mohammedbilalkhan/sagemaker-deployment-guardrail-shadow-testing"}, {"title": "Offline evaluation of machine learning models with noise reduction", "url": "https://image-ppubs.uspto.gov/dirsearch-public/print/downloadPdf/11755947"}, {"title": "Using Feature Flags with Machine Learning Models | ConfigCat Blog", "url": "https://configcat.com/blog/2024/04/23/using-feature-flags-with-ml-models/"}, {"title": "Feature Flags | Statsig Docs", "url": "https://docs.statsig.com/feature-flags/overview"}, {"title": "9 Top Feature Flag Solutions for Modern Product Teams in 2026 | Amplitude", "url": "https://amplitude.com/compare/top-feature-flag-solutions-for-product-teams"}, {"title": "Deployment and Shadow Mode Testing: Validating a New Model on Live Traffic Without User Impact - DYCORA", "url": "https://www.dycora.com/deployment-and-shadow-mode-testing-validating-a-new-model-on-live-traffic-without-user-impact/"}]

I've found comprehensive information on this topic. Here's what the search revealed about these interconnected ML deployment and testing concepts:

Shadow Evaluation & Model Promotion

Shadow testing involves deploying a new model version alongside the current production version, routing a portion of requests to the shadow version, and comparing performance using metrics like latency and error rate to avoid production regressions. Shadow testing allows you to see how new models stack up against your current system by showing what users actually experience?not just what offline tests predict.

Offline Guardrails & No-Regression Checks

Systems can test production/control models against experimental/treatment models by generating experimental requests that represent offline traffic and don't represent unmodified requests from real-world customers. Release gates can be defined before shadow mode begins: acceptable latency, acceptable error rate, acceptable drift range, and acceptable fairness/segment stability thresholds.

Flag-Gated Features & A/B Testing

Feature flags can be configured to augment various processes including Canary Releases, Progressive Deployments, A/B Testing, and can bring these benefits to machine-learning models. In practice, users who see the new feature are treated as the "treatment" group, and users who are gated are treated as the "control".

Integrated Monitoring

Guardrail metrics automatically monitor for negative impacts on feature flag platforms. Shadow mode helps teams detect drift, feature issues, latency problems, and segment-level regressions

before they impact customers, with careful request duplication, isolation, logging, and evaluation gates supporting deployment decisions based on evidence.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

23. user (2026-06-13T20:22:55.616Z)

Web search results for query: "consistency between model prediction and observed evidence refine pseudo-labels temporal action localization verification self-correction"

Links: [{"title": "Pseudo label refining for semi-supervised temporal action localization", "url": "https://www.ncbi.nlm.nih.gov/pmc/articles/PMC11798530/"}, {"title": "Pseudo label refining for semi-supervised temporal action localization | PLOS One", "url": "https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0318418"}, {"title": "Pseudo label refining for semi-supervised temporal action localization - PMC", "url": "https://pmc.ncbi.nlm.nih.gov/articles/PMC11798530/"}, {"title": "Towards Adaptive Pseudo-label Learning for Semi-Supervised Temporal Action Localization", "url": "https://arxiv.org/html/2407.07673v1"}, {"title": "Temporal and Semantic Correlation Network for Weakly-Supervised Temporal Action Localization | ACM Transactions on Multimedia Computing, Communications, and Applications", "url": "https://dl.acm.org/doi/10.1145/3721433"}, {"title": "Learning from Noisy Pseudo Labels for Semi-Supervised ...", "url": "https://openaccess.thecvf.com/content/ICCV2023/papers/Xia_Learning_from_Noisy_Pseudo_Labels_for_Semi-Supervised_Temporal_Action_Localization_ICCV_2023_paper.pdf"}, {"title": "Pseudo-Label Correction Techniques", "url": "https://www.emergentmind.com/topics/pseudo-label-correction"}]

Based on the search results, I found relevant research on this topic. Here are the key findings:

Research on Pseudo-Label Refinement and Temporal Action Localization

A pseudo-label refining (PLR) method based on the teacher-student framework has been proposed, which consists of three key components. The first component is pseudo-label self-refinement which features temporal region interesting pooling to improve the boundary accuracy of TAL pseudo label.

An Instance-level Consistency Discriminator (ICD) has been designed to refine selected pseudo labels on unlabeled videos by removing ambiguous positives and mining potential positives from candidates. This approach focuses on consistency between model predictions and observed evidence.

Self-training methods are achieved by generating pseudo labels and using them to retrain the model, while consistency regularization is achieved by applying different levels of data augmentation and requiring consistent predictions.

Additionally, a temporal consistency loss has been introduced to tackle the location bias problem by penalizing inconsistent boundary predictions from each intrinsic video snippet.

In temporal action localization, context-aware smoothing and online teacher-student correction augment both boundary accuracy and short-action detection.

These approaches collectively address the challenge of verifying consistency between model predictions and observed evidence while refining pseudo-labels for improved temporal action localization.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

24. assistant (2026-06-13T20:23:09.352Z)

This confirms the picture. The "no-regression / shadow / flag-gated promotion" pattern is well-established MLOps (SageMaker shadow testing, feature-flag gating) ? not novel as a primitive,

which is consistent with the survey not claiming it as a research gap. The "consistency between prediction and evidence ? refine" pattern has close prior art (pseudo-label refining, Instance-level Consistency Discriminator, boundary-consistency losses) ? but all of these are *learned, in-loop, silent* refinements, NOT the survey's *report-first, rule-precedence, human-in-the-loop, A/B-gated linter over persisted evidence*. So the survey's claimed differentiation (the linter being report-first and human-gated rather than a learned in-network branch) holds up; no single paper refutes the COMBINATION.

I have enough to render a verdict. Let me compile the structured output.

Summary of findings before the structured call:

- ****All 9 cited papers are REAL.**** No fabrications. arXiv IDs / PMC ID / venues resolve.
- ****Metadata errors to flag (not fabrications):**** (a) LOOCV-bias paper is in ****Science Advances**** (arXiv:2406.01652), not "Briefings in Bioinformatics / GigaScience-class"; (b) SF-Net datasets are THUMOS14/GTEA/BE0ID, not "THUMOS14/ActivityNet"; (c) soccer-HMM title/year is slightly fuzzy but the body of work is real.
- ****Novelty:**** "partial-overlap" is correct and arguably should stay there. The badminton "Point Segmentation" task itself has direct prior art (arXiv:1712.08714) which the survey under-cites in key_papers (it's only alluded to in the rationale). The claimed differentiators are genuine *combinations*, not gaps, and survive refutation searches.

25. user (2026-06-13T20:23:55.246Z)

Structured output provided successfully

26. assistant (2026-06-13T20:24:08.426Z)

Fact-check complete. Summary of what I found:

****Fabrications: none.**** All 9 cited papers are real and resolve to genuine arXiv IDs / PMC ID / venues with correct authorship.

****Metadata errors to fix (not fabrications):****

- The LOOCV-bias paper's venue is wrong ? it's ****Science Advances**** (DOI 10.1126/sciadv.adx6976, arXiv:2406.01652), not the survey's guessed "Briefings in Bioinformatics / GigaScience-class." The PMC ID and the RL00CV technical claim are correct.
- SF-Net evaluates on THUMOS14/GTEA/BE0ID, not "THUMOS14/ActivityNet."
- The soccer-HMM cite is a real body of work but the exact title/year is fuzzy across a cluster of near-identical titles; tighten it.

****Verdict: partial-overlap holds**** (no downgrade to "reinventing"). The adversarial refutation searches surfaced closer prior art the survey under-cites but did not collapse the claimed differentiation:

- ****arXiv:1712.08714**** ("Towards Structured Analysis of Broadcast Badminton Videos") defines our exact task under the established name "Point Segmentation" ? this is the single biggest legibility gap and should be a cited key paper, not just an aside.
- ****Decouple-SSAD (arXiv:1904.07442)**** is the closest prior art for the "detection ? decision" separation.
- ****Pseudo-label-refining / Instance-level Consistency Discriminator (PMC11798530, ICCV 2023)**** is the closest lineage for the reconciler's prediction-vs-evidence loop ? but these are learned/silent, so they don't subsume the report-first human-gated linter; they're just the lineage to cite beyond ASRF.

The survey's defensive evaluation corrections (RL00CV, per-fold variance, tolerance metrics, segmental F1@k, action-context probe) all check out as real and well-targeted; none were oversold. The only "avoid" note: don't claim novelty on the shadow/flag-gated no-regression harness as a primitive ? that's textbook MLOps. Full details are in the StructuredOutput call.