

claude-windows / 985cc286-9e1e-4217-87c3-2d418bc4fa9c

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\985cc286-9e1e-4217-87c3-2d418bc4fa9c\subagents\agent-a4d7a4e88118bba2c.jsonl`
- SHA-256: `37a655c630ce22bf32dd6b9f5f778cf6d9f30dd3b10b4f653ce0014cd292c1ec`
- Source modified: `2026-06-14T09:54:47+00:00`
- Imported at: `2026-07-05T16:48:25+00:00`
- project: `subagents`
- session_id: `985cc286-9e1e-4217-87c3-2d418bc4fa9c`

Transcript

1. user (2026-06-14T09:49:32.070Z)

Collect ALL the golden-set numbers for the Gate-A/B foundational decision and write an implications summary. Owner's framing: NUMBERS INFORM, FOUNDATIONS DECIDE (extensibility, scope of exploration, rigid foundational choices weigh more than raw numbers) ? but make it a data-INFLUENCED choice. In the badminton-highlight-indexer repo (cwd C:/Users/avidu/Projects/badminton-highlight-indexer): READ-ONLY on tracked code; you may RUN evals and WRITE only to scratch/.

Inputs: the 6 persisted golden feature JSONs (output/*_golden_features.json), the experiment harness (backend/eval/experiment.py, backend/eval/fusion_compare.py, backend/eval/fusion_audio.py), backend/eval/eval_stats.py (honest bootstrap CI + paired sign test), and existing baselines (eval_baselines/*.json, docs/QUALITY_ITERATIONS.md). FIRST inspect a sample golden JSON + backend/eval/fusion_golden.py to learn which signals exist as feature COLUMNS (shuttle / person / audio) vs GATES (Gate-A = net_crossings, Gate-B = players_in_court).

RUN a LEAVE-ONE-SIGNAL-OUT ablation on the golden set: full signal set vs leave-one-out for EACH signal group (Gate-A, Gate-B, person, audio, and individual shuttle features where sensible). For each, report held-out ?F1 with bootstrap 95% CI + paired sign test (reuse experiment.compare + eval_stats) = that signal's MARGINAL CONTRIBUTION (does removing it drop quality, quantifiably?). Also the Gate-A specifics: F1 + segment-count/over-seg ratio with Gate-A ON vs OFF (known: ?F1 +0.074, 6/6, p=0.031). If the JSONs lack audio columns, run `python -m backend.eval.fusion\_audio --augment` first.

ASSEMBLE one comprehensive numbers table (these ablations + existing baselines: heuristic/Gen-0/wrapper clean-vs-multicourt, fusion P3 person, P4 audio, Gate-A). Be HONEST about CIs spanning 0 at n=6 (the central caveat).

WRITE scratch/GATE_AB_NUMBERS_AND_IMPLICATIONS.md: (a) the comprehensive numbers table; (b) "Gate-A default flip ? implications" covering numbers AND extensibility / scope-of-exploration / foundational choices (does making Gate-A the default detection constrain or enable future signal-fusion? does it commit a foundational structure that's hard to reverse?); (c) an honest recommendation framing (numbers + foundations, not numbers alone).

GPU + Gemini API are AUTHORIZED if a served-pipeline / end-to-end Gemini run on the golden videos would add real decision value (e.g. end-to-end highlight quality under Gate-A on vs off). BUT the persisted-feature ablation is CPU/fast and decisive for the fusion layer ? assess whether a served Gemini run is WORTH it; do NOT burn a long/expensive Gemini run unless it clearly changes the decision ? FLAG it for the owner with the expected cost/time instead. Run all CPU evals with real commands and put the ACTUAL measured numbers in the doc (no fabrication; if a run errors, report it).

Return: the numbers table + the key implications + the recommendation framing + a yes/no-with-reason on whether a served Gemini run is warranted.

2. assistant (2026-06-14T09:49:37.227Z)

I'll start by reading the session handoff and current state to ramp up, then inspect the golden JSONs and eval harness. Let me pull the repo first and read the key files in parallel.

3. user (2026-06-14T09:49:40.299Z)

Already up to date.

---STATUS---

master

4be16e1 docs(personalization): add PERSONALIZATION_PLAN ? hybrid per-user feature + locked owner decisions (#144)

4. user (2026-06-14T09:49:40.396Z)

output\adarsh_avi_singles_golden_features.json
output\kushagra_singles_golden_features.json
output\largetest_doubles_golden_features.json
output\gbaaddy_golden_features.json
output\testlarge_short_golden_features.json
output\mahadevpura_singles_golden_features.json

5. user (2026-06-14T09:49:40.462Z)

```
1      """Offline golden-set person-feature extractor ? the fusion-P3 measurement's GPU step.
2
3      Mirrors `distill_local`'s CSV-driven flow but ADDS the ~5 person-cue features per
shuttle
4      candidate window, by running `PersonDetector` over the ORIGINAL video. It reuses the
5      EXISTING golden trajectory CSVs (so **no WASB re-run**), needs **no Gemini and no DB**,
and
6      writes one small JSON per video. The LOVO comparison (`fusion_compare.py`) then runs
100% on
7      CPU offline. So the GPU touches exactly one thing ? person detection ? and the lift is
8      measured by the shipped experiment harness; person features enter only as feature
columns the
9      learned scorer weights (no special-casing, default-OFF until the golden ?F1 justifies
them).
10
11      Run (on the GPU box):
12          python -m backend.eval.fusion_golden --manifest output/human_lovo_manifest.json
--out-dir output
13
14      Each output `output/<name>_golden_features.json` is the replayable feature substrate ?
re-run
15      `fusion_compare` with different variants/thresholds forever without re-detecting on the
GPU.
16      """
17      from __future__ import annotations
18
19      import argparse
20      import json
21      import os
22      import time
23      from typing import Any, Dict, List, Optional
24
25      from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
26      from backend.pipeline.detectors import rally_gate
27      from backend.pipeline.detectors import fusion_features as ff
28      from backend.eval import distill_local
29      from backend.eval.calibrate_wasb import load_golden_gts
30      from backend.eval.training import label_candidates
31      from backend.utils.run_telemetry import RunTelemetry
32
33      # The 5 fps/resolution-invariant shuttle features distill_local already computes.
34      SHUTTLE_FEATURE_NAMES = list(distill_local.FEATURE_NAMES)
```

```

35     # The 5 person-cue features (fusion_features.PERSON_FEATURE_NAMES).
36     PERSON_FEATURE_NAMES = list(ff.PERSON_FEATURE_NAMES)
37
38
39     def _derive_video_path(spec: Dict[str, Any]) -> str:
40         """Original MP4 for a manifest entry ? explicit `video_path`, else derived from the
41         golden labels path (`X.rallies.csv` -> `X.MP4`, the golden-set naming
42         convention)."""
43         if spec.get("video_path"):
44             return str(spec["video_path"])
45         labels = str(spec["labels"])
46         if labels.endswith(".rallies.csv"):
47             return labels[: -len(".rallies.csv")] + ".MP4"
48         raise ValueError(f"{spec.get('name')}: no video_path and labels {labels!r} isn't
49         *.rallies.csv")
50
51     def extract_video(spec: Dict[str, Any], detector: Optional[Any] = None,
52                     frame_skip: int = 3, iou: float = 0.5,
53                     log_every_sec: float = 15.0,
54                     telemetry: Optional[Any] = None) -> Dict[str, Any]:
55         """One golden video -> a feature record (shuttle + person features per candidate
56         window).
57         Returns `{name, fps, frame_width,
58         candidates:[{start,end,shuttle},{},person{},label]}, gts}`.
59         `detector` (anything with `detections_over_windows`) is injectable for GPU-free
60         tests;
61         when None a real `PersonDetector` is built (loads torchvision on first use).
62         `telemetry` (a `RunTelemetry` or None) records per-frame velocity + machine-health
63         snapshots
64         from the detection progress callback so a hard kill leaves a forensic runlog.
65         """
66         fps, fw = float(spec["fps"]), float(spec["frame_width"])
67         traj = str(spec["trajectory"])
68         points = TrackNetRunner.parse_trajectory_csv(traj)
69         intervals, x_shuttle = distill_local.build_candidates(traj, fps, fw)
70         gts = load_golden_gts(str(spec["labels"]))
71         labels = label_candidates(intervals, gts, iou)
72         # Net axis for the person two-sided split ? reuse rally_gate's camera-agnostic
73         estimate
74         # (parity with the shuttle net-crossing gate) rather than hard-coding an
75         orientation.
76         axis, net_px, _span = rally_gate.estimate_play_axis(points)
77         video_path = _derive_video_path(spec)
78
79         if detector is None:
80             from backend.pipeline.detectors.person_detector import PersonDetector
81             detector = PersonDetector({"indexing": {"use_gpu": True}})
82
83         # Single-pass person detection over ALL candidate windows (open the video ONCE, read
84         # forward ? no per-window seek). Then compute features per window (pure, no GPU).
85         n = len(intervals)
86         win_frames = [(round(s * fps), round(e * fps)) for (s, e) in intervals]
87         print(f"[fusion_golden] {spec['name']}: {n} windows ? person-detecting in one
88         pass...", flush=True)
89         t0 = time.monotonic()
90         state = {"last": t0}
91
92         def _progress(done: int, total: int) -> None:
93             if telemetry is not None:
94                 telemetry.snapshot(done, total, phase=spec["name"])
95             now = time.monotonic()
96             if now - state["last"] >= log_every_sec or done >= total:
97                 el = now - t0

```

```

91         rate = done / el if el > 0 else 0.0
92         eta = (total - done) / rate if rate > 0 else 0.0
93         pct = 100 * done // total if total else 0
94         print(f"[fusion_golden]    {spec['name']}: detect frame {done}/{total}
({pct}%) "
95             f"| {el:.0f}s elapsed | ETA {eta:.0f}s", flush=True)
96         state["last"] = now
97
98     per_window = detector.detections_over_windows(video_path, win_frames, frame_skip,
99                                                  progress_cb=_progress)
100
101     candidates: List[Dict[str, Any]] = []
102     for i, ((s, e), xs) in enumerate(zip(intervals, x_shuttle)):
103         det = per_window[i]
104         det_fw = det.get("frame_width") or fw
105         det_fh = det.get("frame_height") or 0.0
106         # Normalise the (pixel) net coord to 0-1 on the play axis, matching the foot-
point
107         # normalisation fusion_features uses (x/width, y/height).
108         if axis == "x":
109             net_norm = net_px / det_fw if det_fw > 0 else 0.5
110         else:
111             net_norm = net_px / det_fh if det_fh > 0 else 0.5
112         sf, ef = win_frames[i]
113         shuttle_in = [p for p in points if sf <= p.frame <= ef]
114         person = ff.compute_person_features(det["frames"], shuttle_in, det_fw, det_fh,
115                                            court_polygon=None, net=(axis, net_norm))
116         candidates.append({
117             "start": float(s), "end": float(e),
118             "shuttle": {k: float(v) for k, v in zip(SHUTTLE_FEATURE_NAMES, xs)},
119             "person": {k: float(v) for k, v in person.items()},
120             "label": int(labels[i]),
121         })
122     return {
123         "name": spec["name"], "fps": fps, "frame_width": fw,
124         "candidates": candidates,
125         "gts": [[float(a), float(b)] for a, b in gts],
126     }
127
128
129 def run(manifest_path: str, out_dir: str, frame_skip: int = 3, iou: float = 0.5,
130        fresh: bool = False, smallest_first: bool = False) -> List[str]:
131     with open(manifest_path, encoding="utf-8") as f:
132         specs = json.load(f)
133     if smallest_first:
134         # Process small videos first for quick feedback (the slow full match lands
last);
135         # the LOVO result is order-independent. Proxy size by the trajectory CSV (one
row/frame).
136         specs = sorted(specs, key=lambda s: (os.path.getsize(s["trajectory"])
137                                             if os.path.isfile(s.get("trajectory", ""))
else 0))
138     os.makedirs(out_dir, exist_ok=True)
139     # Run-telemetry black box: a long GPU run can be OS-killed with no traceback, so
keep a tiny
140     # rotating runlog of velocity + machine health for the post-mortem. Constructing it
must never
141     # break extraction, so guard the build and fall back to no telemetry on any error.
142     telem: Optional[RunTelemetry]
143     try:
144         telem = RunTelemetry(os.path.join(out_dir, "run_telemetry.jsonl"),
label="fusion_golden")
145     except Exception: # noqa: BLE001 - telemetry wiring must never break the run it
observes
146         telem = None

```

```

147     written: List[str] = []
148     for vi, spec in enumerate(specs, 1):
149         out = os.path.join(out_dir, f"{spec['name']}_golden_features.json")
150         # Resumable: skip a video whose JSON already exists (a re-run picks up where a
prior
151         # run left off instead of re-detecting). --fresh forces a redo.
152         if not fresh and os.path.isfile(out):
153             print(f"[fusion_golden] ({vi}/{len(specs)}) {spec['name']}: already
extracted "
154                 f"-> {out} (use --fresh to redo)", flush=True)
155             written.append(out)
156             continue
157         print(f"[fusion_golden] ({vi}/{len(specs)}) {spec['name']}: extracting...",
flush=True)
158         t0 = time.monotonic()
159         if telem is not None:
160             telem.event("video_start", spec["name"])
161         try:
162             data = extract_video(spec, frame_skip=frame_skip, iou=iou, telemetry=telem)
163         except Exception as e: # noqa: BLE001 - record machine state at death, then re-
raise
164             if telem is not None:
165                 telem.event("error", spec["name"], err=str(e))
166             raise
167         n_pos = sum(c["label"] for c in data["candidates"])
168         with open(out, "w", encoding="utf-8") as f:
169             json.dump(data, f, indent=2)
170         if telem is not None:
171             telem.event("video_done", spec["name"],
172                         windows=len(data["candidates"]), positives=n_pos)
173         print(f"[fusion_golden] ({vi}/{len(specs)}) {spec['name']}: DONE in
{time.monotonic() - t0:.0f}s "
174             f"| {len(data['candidates'])} windows ({n_pos} positive) |
{len(data['gts'])} golden rallies "
175             f"-> {out}", flush=True)
176         written.append(out)
177     return written
178
179
180     def main() -> None:
181         ap = argparse.ArgumentParser(description="Extract golden-set shuttle+person features
(fusion P3 GPU step).")
182         ap.add_argument("--manifest", default="output/human_lovo_manifest.json")
183         ap.add_argument("--out-dir", default="output")
184         ap.add_argument("--frame-skip", type=int, default=3)
185         ap.add_argument("--iou", type=float, default=0.5)
186         ap.add_argument("--fresh", action="store_true", help="re-extract even if a JSON
already exists")
187         ap.add_argument("--smallest-first", action="store_true",
188                         help="process smaller videos first for quick feedback (LOVO is
order-independent)")
189         args = ap.parse_args()
190         run(args.manifest, args.out_dir, args.frame_skip, args.iou, fresh=args.fresh,
191             smallest_first=args.smallest_first)
192
193
194     if __name__ == "__main__":
195         main()
196

```

6. user (2026-06-14T09:49:40.742Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-452 of 1835 total (86106 tokens, cap 25000). Call Read with offset=453 limit=452 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-

reminder>

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (FUSION P3/P4 MEASURED + extraction COMPLETE + telemetry #134
merged + Cooler Boost). **THE NUMBERS
13     (honest LOVO n=6, IoU?0.5, C1 bootstrap-CI + sign-test):** shuttle-only F1=**0.307**;  

**+person (P3) ?F1=?0.021** (?6.7%, CI
14     [?0.165,+0.162], sign 2W/4L p=0.688, **CI-clears-0: NO**) ? NO lift, stays default-OFF
(no court polygons ? `players_in_court`
15     counts spectators; per-video wash: testlarge +0.40 / gbaaddy ?0.22). **+person+audio (P4
incremental) ?F1=+0.079** (+27.7%, F1
16     0.286?+0.366**, CI [?0.054,+0.203], sign 4W/2L p=0.688, **CI-clears-0: NO**) ?
PROMISING (wins 4/6, largetest/mahadevpura
17     +0.27; endpoint 0.366 > shuttle-only 0.307) but n=6 can't confirm. **VERDICT: neither
promotable (both CIs span 0); audio is
18     the lead ? need more matches + measure audio on shuttle-only directly (drop the person
handicap).** Recorded in
19     `eval_baselines/experiment_leaderboard.json`; honest negative/uncertain results KEPT
([[paper-coauthor-aim]]). **LANDED THIS
20     SESSION:** result logged to `docs/QUALITY_ITERATIONS.md` (#137);
`experiment_leaderboard.json` tracked (#139); **weak-signals
21     audit ?** (person/audio/court-bounds all present, default-OFF, no-special-casing,
retained ? refuted=false ? [[weak-signals-retained]]);
22     7-page paper-style **quality report PDF** at
`OneDrive\Documents\Rally_Detection_Quality_Report.pdf` (gen `scratch/quality_report.py`;
23     conservative academic restyle, strict-reviewer-approved); `FusionGoldenExtract`
scheduled task **unregistered** (cleanup); tree
24     clean on `master`. ?? **NEXT (fusion track):** (1) **wire `run_telemetry` (#134) into
fusion_golden + the ps1 wrapper** (deferred,
25     ready, mergeable); (2) **measure audio on shuttle-only directly** (quick eval, reuses
the 6 `output/*_golden_features.json`, NO new
26     data ? audio's standalone lift without the person handicap); (3) **grow the golden
corpus via Roora** (owner's job ? gives power
27     to confirm the +0.079 audio lead) + **activate court polygons** (re-measure person
properly ? the ?0.021 is partly a
28     `court_polygon=None` artifact); (4) future **shuttle-dropped** cue ([[weak-signals-
retained]]). **OTHER track (per
29     [[session-handoff]] ? the stated IMMEDIATE PICKUP): EXECUTION-POLICY P3 then P4**
(owner-gated; #132/#133/#135 done). **Extraction infra** (all working): detached Scheduled Task
`FusionGoldenExtract`
30     runs `scratch/extract_until_done.ps1` (self-restart loop + in-process keep-awake + GPU-
busy gate: waits if free VRAM<3500MB) ?
31     deaths were `0xC000013A` console-control kills (NOT crash/sleep/TDR/OOM ? verified via
event log), auto-recover via 3-min
32     repetition trigger + restart-on-failure. **Cooler Boost** (MSI Creator Center) dropped
85?67°C, SM ~1100?~1300MHz, velocity
33     ~26?~50 fps (the rest is Max-Q power cap, not thermal).
34
35     **Checkpoint:** 2026-06-14 (EXECUTION-POLICY framework: design + P1 + P2 SHIPPED ? **PR
#132 (design) + #133 (P1) + #135 (P2)
36     MERGED**, self-merged on green CI). Owner asked to harden the GX010125 Gemini hang as a
**policy-driven, per-job,
37     self-scheduling** system (default WAIT_AND_RESUME, ABORT, POLL EVERY_N_SEC?, **no master
node**) with **very high coverage**
```

```

38      (failures cause delays) and **clean poll-logging** (visible but not polluting). ***132**
= design doc `docs/EXECUTION_POLICY.md`
39      (validated the owner's architecture: per-job policy = data on the job; self-schedule via
re-enqueue; no master ? rides the
40      existing single-worker executor + `jobs` table + crash sweep; DB is the coordinator).
***133 (P1, the actual fix)** =
41      `backend/infrastructure/execution_policy.py`: `ExecutionPolicy` + `run_bounded` (HARD
`op_timeout_sec` per-attempt ? the
42      hang-killer; retries failures AND hangs; hung attempt's daemon thread abandoned) +
`poll_until` (bounded poll loop). Clocks
43      injected ? **12 tests, 100% module coverage, instant**. Refactored `gemini.py`:
`generate_content`? `run_bounded` (a
44      never-returning generate previously blocked FOREVER = the GX010125 hang; now bounded),
processing-wait? `poll_until` (**per-poll
45      logs at DEBUG, not an INFO line every 10s** ? the no-pollution ask). 23 gemini tests
stay green. ***135 (P2, substrate)** =
46      DB migration **v3** (`jobs.policy` JSON + `next_poll_at`; existing rows survive via
ALTER) + `set_job_waiting(delay_sec)`
47      (self-schedule: park 'waiting' until due) + **`claim_due_job()`** (atomic compare-and-
set claim of queued/due-waiting ?
48      running; concurrent-safe, no master) + `create_job(policy=)`/`get_job` policy.
**Additive, no executor change**; 10 tests;
49      version asserts bumped 2?3. ? **THIS IS THE OWNER'S REAL GPU MACHINE**
(torch+CUDA/cv2/WSL/Gemini live). ?? **REMAINING
50      (owner-gated, NOT done ? deliberately not rushed at session tail given the high-coverage
bar): P3** = wire the executor
51      (main.py) to a `claim_due_job` self-scheduling loop honoring WAIT_AND_RESUME
(release+reschedule) + **resumable AI handoff**
52      (persist per-window confirmation; resume only unconfirmed ? the GX010125 reel becomes
exhaustive across resumes, no wasted
53      Gemini spend); **P4** = apply the policy to the WSL/WASB GPU pass + ffmpeg. Each = its
own distinct PR. Ties [[paper-coauthor-aim]] indirectly (infra).
54
55      **Checkpoint:** 2026-06-14 (ACTED ON FIRST-A/B FINDINGS: Gate-A PROMOTED + learned FIXED
+ GX010125 run ? **PR #130 + #131
56      MERGED** to master `2e5892e?`; self-merged on green CI). ? **THIS IS THE OWNER'S REAL
GPU MACHINE** (torch+CUDA, cv2, WSL
57      Ubuntu, Gemini key in .env all present) ? not the dev container the older CLAUDE.md note
assumes. ***130 = learned-variant
58      FIX**: additive/default-OFF `Variant.tune_threshold` (operating point on the TRAIN fold)
+ `calibrate` (Platt/isotonic, C2);
59      validated on n=6 golden ? **mean held-out F1 0.307?0.423 (+0.116), testlarge 0.000?0.300
(un-zeroed), now BEATS CONTROL
60      0.388**; threshold-in-LOVO is the big lever. ***131 = Gate-A PROMOTED** (first cue
promoted via the harness; doc trail in
61      QUALITY_ITERATIONS ? measured win dF1 +0.074, 6/6, p=0.031; recommended
`indexing.fusion.min_crossings=2`; **live served-default
62      flip stays OWNER-GATED** until a real fusion-path run confirms ? fusion P2 is unverified
on real video) **+ new
63      `docs/PER_VIDEO_OUTPUT_LAYOUT.md`** (one dir per video; fixes the hardcoded `"output"`
joins in `trajectory_hybrid.py:~176`
64      so `--output-dir` truly isolates). **GX010125_1080p RUN (offline, no GPU):** ? key
finding ? `wasb_runner.run_predict` has
65      **NO cache-hit early-return**, so the full pipeline would RE-RUN WASB (~tens of min on
1080p) + overwrite the cached CSV; the
66      cached trajectory is reusable ONLY via offline re-score. So ran
`scratch/run_gx010125_rescore.py` = parse cached
67      `output/wasb/GX010125_1080p_89605e80ed01_wasb.csv` ? 45 candidates ? CONTROL 45 rallies
(11.0min) vs **Gate-A 27** (9.3min),
68      Gate-A **dropped 18 low-crossing windows (~1.7min)** of likely held-shuttle/feed FPs ?
isolated `output/GX010125_run/` (no
69      pollution; no golden labels for this footage ? divergence not F1). A FULL Gemini-
confirmed highlight reel = a separate
70      ~tens-of-min GPU+API run (owner kicks off). Repro:
`scratch/{run_gx010125_rescore,ab_n6_lovo,ab_eval_demo}.py`. All in

```

```

71     isolated worktrees off master (siblings concurrent). ?? NEXT (owner-gated): real fusion-
path run to confirm Gate-A served-side
72     + flip the live default; impl `PER_VIDEO_OUTPUT_LAYOUT` L1; Tier-2/3 §13 corrections;
analyzers/reconcilers A1?A5/R1?R3. Ties [[paper-coauthor-aim]].
73
74     **Checkpoint:** 2026-06-14 (FIRST GOLDEN A/B EXPERIMENT + HARNESS HONEST-REPORTING WIRED
? **PR #128 + #129 MERGED**
75     to master `f680b8d`, self-merged on green CI per [[working-agreement-merging]]; owner-
directed momentum). Ran the
76     **first real end-to-end A/B on the n=6 golden corpus, 100% OFFLINE** (re-scored
persisted WASB trajectories via
77     `distill_local.build_candidates`, no GPU ? the "persist once re-score" thesis proven).
One candidate set re-scored
78     three ways: **Gate-A (net_crossings?2) is a MEASURED WIN ? ?F1 +0.074, 6/6 matches,
sign-test p=0.031, CI
79     [+0.042,+0.104], seg-count-ratio 1.68?1.01** (over-seg nearly eliminated); the **window-
level 5-feature logistic did
80     NOT clear the bar** (?0.081, 2W/4L; testlarge zeroed = threshold/calibration artifact ?
C2 motivation; **distinct from
81     the stronger per-frame `rally_seq_proto` 0.583** ? don't conflate). Headline: **the
framework held up ? C1 promoted one
82     variant and refused another (right call both ways); C4 surfaced over-seg tIoU hid.**
**#128** = archived posterity doc
83     `docs/archives/quality-iterations/2026-06-first-ab-experiment/` (full tables) +
archives-index + live QUALITY_ITERATIONS
84     pointer. **#129 = (b) wired C1+C4 into the CORE `experiment.compare()`**
(additive/default-ON `honest_stats=True`:
85     `honest` block = per-metric bootstrap CIs + paired ?F1 CI + exact sign test + F1@k +
segment_count_ratio; `record_experiment`
86     carries honest ?F1; existing fields byte-unchanged, `honest_stats=False` = legacy
shape). Reuses #124 helpers; complements
87     sibling **#127** which wired the same C1 into `fusion_compare` (no dup). 19 existing
experiment tests pass + 3 new;
88     ruff+mypy+CI all green. Repro `scratch/ab_n6_lovo.py` + `scratch/ab_eval_demo.py`
(machine-side; read gitignored
89     manifest/external golden). All done in **isolated worktrees off master** (siblings
merged #121-#127 concurrently). ?? NEXT
90     (owner-gated): promote Gate-A via the no-regression gate (it's a real measured win); fix
the logistic (C2 calibration +
91     threshold-inside-LOVO); Tier-2/3 §13 corrections; analyzers/reconcilers impl A1?A5/R1?R3
(none started). Ties [[paper-coauthor-aim]].
92
93     **Checkpoint:** 2026-06-14 (FUSION P3/P4 MEASUREMENT BUILT + GPU EXTRACTION RUNNING ?
master `d7b1f51`; suite 514 green).
94     The "modeling" track for the person-cue (P3) + audio (P4) fusion lift ? all default-OFF,
no-special-casing, harness-measured.
95     **Merged this run (mine): #121** fusion P3 (`backend/eval/fusion_golden.py` offline
person-feature EXTRACTOR + `fusion_compare.py`
96     LOVO driver ? reuses build_candidates/PersonDetector/fusion_features/experiment.compare,
NO new scoring); **#122** P4
97     (`backend/pipeline/detectors/audio_features.py` + `backend/eval/fusion_audio.py`
incremental audio cue); **#123** progress-logging
98     + resumability (`--fresh`); **#125** single-pass
`PersonDetector.detections_over_windows` (fixed the per-window-peek O(N)
99     pathology ? ~28s/window ? one sequential decode); **#126** `--smallest-first`; **#127**
honest ?F1 = bootstrap CI + paired
100     sign test on per-video held-out F1 (wires #124's `eval_stats.paired_delta_ci`; "never a
bare mean"). Concurrent sessions
101     merged #119 (analyzers framework), #120 (run.log), #124 (eval-honesty utils:
eval_stats/score_calibration/segmentation_metrics).
102     **? GPU EXTRACTION via DETACHED SCHEDULED TASK `FusionGoldenExtract`** (runs
`scratch/extract_until_done.ps1` ? Task Scheduler
103     owns it, OUTSIDE Claude's job object ? survives turn/context/session boundaries; in-
process keep-awake + self-restart loop until
104     6/6; clean UTF-8 logs to `output/fusion_golden.log` + `scratch/extract_wrapper.log`).

```



```

Loops `python -u -m backend.eval.fusion_golden
105     --manifest output/human_lovo_manifest.json --out-dir output --smallest-first` ? 6
`output/<name>_golden_features.json` (GITIGNORED
106     + RESUMABLE: skips any video whose JSON exists). **3/6 DONE** (testlarge, mahadevpura,
largetest); remaining
107     kushagra?gbaaddy?adarsh(last, 84k frames). **MONITOR:** JSON count + `tail
output/fusion_golden.log` + `nvidia-smi`;
108     `Get-ScheduledTask FusionGoldenExtract|Select State`. **RESTART (resumes, skips done):**
`Start-ScheduledTask -TaskName
109     FusionGoldenExtract`. **STOP:** `Stop-ScheduledTask -TaskName FusionGoldenExtract` +
`Get-Process python|Stop-Process -Force`.
110     ?? **WHY DETACHED (debugged 2026-06-14):** Claude-bg-task runs DIED TWICE silently
(largetest 61%, kushagra 6% @137s), NO Python
111     traceback AND ? verified via Windows event log ? NO sleep (Kernel-Power 42), NO TDR
(Display 4101), NO WER crash, NO RADAR-OOM
112     (no event 1003; 18GB free). Traceless vanish = external `TerminateProcess`; common
factor = both were Claude-managed bg tasks
113     (job-object kill-on-close at a context boundary fits death #1 @ the prior context
summary). Scheduled Task = immune; GPU access
114     CONFIRMED in the interactive task session (53?64% util, 3.1GB). GPU util while running =
20?64% bursty ? batch=1 Faster-RCNN
115     latency-bound (small kernels + CPU NMS + per-frame ByteTrack), NORMAL, model IS on CUDA,
NOT decode-bound; batching the forward
116     pass is the next-extraction speed lever (changes features, so NOT mid-run). ~1.5-2h
total. Stale Claude bg keep-awakes stopped. **DEATH #3 + GPU-job check (2026-06-14):** died a
3RD time (gbaaddy 69%)
117     ? `LastTaskResult=0xC000013A` = STATUS_CONTROL_C_EXIT = terminated by a **console-
control event / CTRL+C** (NOT crash/sleep/TDR/OOM
118     ? explains the no-trace vanish; killed the detached task too, so NOT Claude's job
object). GPU-contention angle CHECKED: **NO
119     competing GPU compute job** (only desktop apps + our 1 python; 5GB free; HW is a
**laptop** RTX 2070S **Max-Q**, MSI). Mitigations:
120     (1) task re-registered with **3-min repetition trigger + restart-on-failure (count 10,
1-min)** ? auto-recovers within minutes of
121     ANY death; (2) **GPU-busy gate in the ps1** (owner policy = *wait until free, auto-
resume*): before each python launch, if free
122     VRAM <3500MB (another job has the GPU) it waits 30s + re-checks, then launches. **??
WHEN IT COMPLETES ? the remaining steps for
123     the NUMBERS:** (1) `python -m backend.eval.fusion_compare --features-dir output
--record` ? honest P3 LOVO ?F1 (shuttle-only
124     vs +person) w/ CI+sign-test; (2) `python -m backend.eval.fusion_audio --augment
--compare --record` ? P4 incremental audio
125     ?F1; (3) log both to `docs/QUALITY_ITERATIONS.md` I3 (HONEST measured numbers only ?
[[paper-coauthor-aim]]). Honest caveat:
126     NO court polygons ? `players_in_court` counts spectators (3-7 on "single-court"
Mahadevpura); robust signals = motion +
127     `shuttle_player_colocation`. GPU = RTX 2070S, person detection on Windows torch cu124
(NO WSL). Branch protocol held through
128     heavy concurrent merges (#119-127) ? [[gotcha-shared-checkout-branch-collisions]]. On
master, clean.
129     **STORAGE (this session, owner-approved "delete cache + compact WSL"):** deleted the
stale **93GB GX010125 frame cache**
130     (WSL `~/clips` 95?2.3GB; WSL fs used 152?60GB, 896GB free; C: 53?72GB free) ? kept the
2.2GB GX010125 proxy. The WSL
131     `ext4.vhdx` was `--set-sparse true` + `fstrim`'d (946GB trimmed) but this WSL build
(2.3.24) did NOT auto-deallocate on
132     shutdown ? the file is still **~221GB allocated on C:**. **DEFERRED to post-extraction:
the physical ~150GB C: reclaim via
133     `wsl --shutdown` then elevated `diskpart` ? `select vdisk
file="...\CanonicalGroupLimited.Ubuntu_79rhkp1fndgsc\LocalState\ext4.vhdx"`
134     ? `compact vdisk`** (heavy C: I/O would fight the GPU run's video reads). [[gotcha-long-
runs-gpu-wsl]] · [[cache-hygiene-policy]].
135
136     **Checkpoint:** 2026-06-14 (RALLY-DETECTION FRAMEWORK doc + LITERATURE REVIEW + eval-
honesty code ? **PR #119 + #124 MERGED**

```

```

137     to master via self-merge on green CI per [[working-agreement-merging]]; owner-directed).
**PR #119** (`docs/ANALYZERS_AND_RECONCILERS.md`):
138     reframed the analyzers/reconcilers design into a **multi-scope signal-fusion framework**
(frame detectors / window analyzers /
139     **session analyzers**); warm-up `SessionAnalyzer` + shuttle-state frame-detector folded
in) + **§13 = a verified multi-agent
140     literature-review** (8 patterns surveyed, ~40 citations adversarially checked, **ZERO
fabrications**; author-line fixes applied).
141     **Verdict (honest):** the skeleton is STANDARD prior art ? candidate-windows+scorer =
Temporal Action Localization;
142     reconciler-as-refinement = Temporal Action Segmentation post-proc; "signal not verdict"
= late/score-level fusion + stacked
143     generalization (~20 yrs old, already in racket sports). **Only the report-first
reconciler-as-linter bundle rated
144     `novel-combination`**; warm-up-as-soft-feature = thinnest gap. **Cite, don't claim**;
terminology realigned (producer?labeling
145     function, scorer?stacked meta-learner, LOVO?leave-one-group-out, tIoU?F1@k/mAP). Ties
[[paper-coauthor-aim]] (§13 is the honest
146     related-work backbone). **PR #124** implements the **Tier-1 measurement-honesty course
corrections C1?C4** as PURE ADDITIVE /
147     default-OFF helpers (no existing compare/LOVO/metrics behaviour changed):
`backend/eval/eval_stats.py` (C1 `mean_with_ci`/
148     `bootstrap_ci`/`paired_sign_test`/`paired_delta_ci` ? never a bare LOVO mean at n?6; C3
`out_of_fold_predictions`/`grouped_folds`
149     ? no stacking leakage), `score_calibration.py` (C2 Platt+isotonic+brier/ECE ? comparable
cue confidences), `segmentation_metrics.py`
150     (C4 F1@{10,25,50} + mAP@tIoU + segment_count_ratio ? make over-segmentation tIoU is
blind to visible). 24 tests; ruff+mypy clean;
151     all 4 CI lanes green (full-suite coverage floor held). Both PRs done in **isolated git
worktrees** (off master) to dodge the
152     recurring shared-checkout branch-flips (siblings merged #117 fusion-P2, #118, #120, #123
concurrently; **#116 closed ? its content
153     had landed via #117**). ?? NEXT (owner-gated): **wire C1?C4 into the DEFAULT harness
reporting** (`compare`/`run_regression.py`
154     emit CIs + F1@k/edit-score instead of a bare mean ? changes reported numbers, so its own
PR) + the Tier-2/3 corrections in §13
155     (global reconciler decode, cue-correlation modeling, temporal smoothing, external
ShuttleSet eval). Analyzers/reconcilers IMPL
156     (A1?A5/R1?R3) still owner-gated, none started.
157
158     **Checkpoint:** 2026-06-14 (RUN.LOG READABILITY for manual QA ? **PR #120 MERGED ?
master `6076eef`**); owner asked ? self-merged
159     on verified-green CI per [[working-agreement-merging]] (squash; remote+local branch
deleted); was branch `feat/runlog-readability`
160     off `35a50f3`; CODE, minimal/config-gated). Owner's manual-QA loop = run the tool on a
video ?
161     eyeball the rally CSV against `run.log`. Two readability changes: (1) **suppress HTTP
chatter** ? new central helper
162     `backend/utils/logging_setup.py::setup_logging(verbose_http=False)` raises
`httpx/httpcore/urllib3/google_genai/google_genai/
163     google.auth/grpc` to WARNING (NOTSET to restore), gated by NEW knob
`logging.verbose_http: false` (`LoggingConfig` on
164     `AppConfig`); our `backend.*` loggers untouched; wired into `run_process_and_stitch.py`
(at import + re-applied
165     post-config-load). (2) **end-of-run rally summary** ? new pure helper
`backend/tools/export.py::format_rally_summary` (next to
166     `segments_to_csv`/`basic_stats`): header (total rallies + total play HH:MM:SS) + table
`# start end dur(s) shot_count` from the
167     produced play_segments so log?CSV line up row-for-row; handles metadata dict-or-JSON-
string, missing shot_count?`-`; HH:MM:SS
168     via local `_fmt_hms` matching the segmenters' `_fmt_ts`; logged as ONE block at end of
`main()` (also on the all-filtered
169     early-return). Tests: `test_logging_setup.py`, `test_rally_summary.py`, + `verbose_http`
default/override in `test_config.py`.
170     Suite **477 green**, ruff (`.`) + mypy (`backend training`) clean. ? Commit-msg first

```

```

attempt mangled (PowerShell `@'?'@`
171     here-string in the *Bash* tool ? literal `@`; amended via `-F file`) ? Bash tool ?
PowerShell. **BRANCH-COLLISION CLEANUP
172     (owner asked):** after merge, LOCAL checkout left clean ? on
`master`==`origin/master`==`6076eef`, 0 ahead/behind; pruned 6 stale
173     LOCAL branches mapped to MERGED/CLOSED PRs (#92/#104/#111/#106/#102 + #116-closed-
redundant) + the stale remote-tracking refs;
174     **kept** `docs/analyzers-framework` (#119 OPEN). Working tree = only the intentional
untracked carryovers (`scratch/`,
175     `artifacts/`). REMOTE still has 15 prunable branches (14 merged-PR + #116) ? **offered
to prune, owner did NOT answer the
176     question ? left untouched** (must KEEP `feat/khelacharya-shot-intelligence` #11 =
deliberately-retained sibling layer per
177     [[sibling-internship-repo]], + #119). ?? NEXT: owner does the machine-side check (run on
a real video ? eyeball the quieter log +
178     summary block; dev container can't run the GPU pipeline); remote-branch prune available
on request.
179
180     **Checkpoint:** 2026-06-14 (ANALYZERS DOC ? FRAMEWORK + paper framing ? **PR #119 OPEN,
owner-gated, do NOT merge**); branch
181     `docs/analyzers-framework` off master `35a50f3`; docs-only). Owner fed two more ideas
this session and asked to fold them in
182     and treat the doc as a **reusable framework / research paper**, not a project patch
(ties [[paper-coauthor-aim]]). Extended
183     the on-master `docs/ANALYZERS_AND_RECONCILERS.md` (which had landed via #117) from "two
seams" into a **multi-scope
184     signal-fusion framework**: producers emit SIGNALS at **3 temporal scopes** ? frame
detectors (`RallyCue`) / window analyzers
185     / **session analyzers** ? all consumed as feature columns by the learned scorer (spine
unchanged: **no special-casing**, signal
186     ?scorer, never `if/else`). Folded in: (1) **warm-up/practice** as a new
**`SessionAnalyzer`** tier (whole-timeline span ? an
187     `in_warmup` membership feature the scorer DOWN-WEIGHTS, NOT a hard `if t<warmup_end:
drop`; measurable on golden TODAY since
188     warm-up is already labeled negative); (2) **shuttle-state**
(`in_air`/`racket_contact`/`inactive`) placed correctly as a
189     **frame-scope detector** (a MULTI_SIGNAL_FUSION_PLAN `RallyCue`, NOT a window analyzer)
? derive-from-existing-signals first
190     (velocity+visibility+direction-reversals+audio thwack), dedicated per-frame model ONLY
where WASB trajectory gaps bite, with
191     the per-frame-label blocker flagged. Added **$10 generality & research positioning**
(weakly-supervised temporal segmentation;
192     standard-vs-contributed) with an explicit **attribution honesty caveat**: NO fabricated
citations/results, related-work +
193     ablations + >1-sport still owed, n=6 = direction-check only ? the honest backbone the
paper aim needs. Diff = `2-seam?framework`
194     (310+/137?). ? Shared checkout had flipped to master mid-task (the sibling's collision
lesson recurred) ? re-verified branch,
195     branched fresh, normal commit (no in-flight WIP left to isolate; #117/#118 already
merged so all cross-links resolve). PR #116
196     stays CLOSED (content on master via #117). ?? NEXT (owner-gated): review #119; impl
tracks unchanged (fusion P3 LOVO, then
197     analyzers/reconcilers A1?A5 / R1?R3 per the doc's phasing ? each ONE PR, default OFF,
promote on measured LOVO lift).
198
199     **Checkpoint:** 2026-06-14 (FUSION P2 SHIPPED + collision cleanup + paper aim ? master
`35a50f3`; suite 467 green; 0 open PRs).
200     **PR #117 MERGED = `FusionSegmenter`** (fusion P2, registry `fusion_hybrid`, **default-
OFF**): shuttle-primary (windows seed
201     from WASB/TrackNet) + adjudicating cues ? `net_crossings` (Gate A, always persisted,
GPU-free via rally_gate w/ a reused
202     axis estimate) + optional person cue (Gate B) persisting the ~5 `fusion_features.py`
features (pure-tested) via
203     `PersonDetector.detections_per_frame` over each window's ORIGINAL-video frame range.
Plugs in via **two IDENTITY hooks** on

```

```

204     `trajectory_hybrid` (`_enrich_candidate_stats`/_`select_for_handoff`) ? wasb/tracknet
byte-identical (adversarially verified).
205     Served gating optional+default-OFF (`fusion.min_crossings`/_`min_players`) ? gates the
handoff while persisting the full
206     SUPERSET (offline substrate intact). Court mask = optional `fusion.court_polygon`.
**Design-first** (`fusion-p2-design`
207     workflow) ? **31-agent adversarial review** (25 findings?19 confirmed; headline = NO-
REGRESSION verified) ? fixes folded
208     (frame-dim guards, round() frame idx, rally_gate `estimate=` reuse, honest torch
docstrings, +6 tests). **PR #118 MERGED** =
209     README index fix (QUALITY_ITERATIONS + EXPERIMENT_HARNESS P1-SHIPED). **? SHARED-
CHECKOUT COLLISION (lesson):** the
210     spawn_task DESIGN session ran in THIS checkout, committed `0f673ff` (the analyzers
design doc) + flipped the branch BETWEEN
211     my branch-point check and my `checkout -b` ? my fusion branch rode on top ? **#117's
squash absorbed the design doc**, so
212     **`docs/ANALYZERS_AND_RECONCILERS.md` is on master via #117** and **PR #116 was CLOSED
as redundant** (content intact on
213     master; #116's stale "OPEN do-not-merge" checkpoint below is superseded). Re-verify `git
branch` AND that no sibling commit
214     landed since, before `checkout -b`. **Owner side-aim ratified: co-author a paper**
"local, cheap & efficient rally detection
215     via multiple clever ideas" ? [[paper-coauthor-aim]]; `docs/QUALITY_ITERATIONS.md` is the
ablation-ready empirical backbone
216     (honest measured results only). **Analyzers+reconcilers DESIGN on master (owner-gated,
NOT started):** injectable per-window
217     analyzers (e.g. net-hit service-fault) + a report-first reconciliation pass; spine =
signals?learned-scorer, **no
218     special-casing** (owner directive). ?? NEXT (owner-gated): **fusion P3** (person
features ? `experiment.compare` LOVO lift on
219     the 6 golden videos ? the eager-person-compute cleanup lands here) + measure the Gate-A
held-shuttle drop on GPU; then
220     analyzers/reconcilers impl; P4 audio. Earlier this session: PR #114 harness, #115
person-cue extraction (see checkpoints below).
221
222     **Checkpoint:** 2026-06-14 (DESIGN DOC: two extensibility seams ? **PR #116 OPEN, owner-
gated, do NOT merge**; docs-only,
223     branch `docs/analyzers-and-reconcilers` off master `df7d5df`). Produced
`docs/ANALYZERS_AND_RECONCILERS.md` (367 lines, a
224     peer to EXPERIMENT_HARNESS + MULTI_SIGNAL_FUSION_PLAN) + a 1-line `docs/README.md`
cross-link. **Spine (owner directive):
225     NO special-casing in the decision path** ? every analyzer/reconciler emits a SIGNAL
(value+confidence); the SCORING MODEL
226     learns to weight it (already true in miniature: `experiment.predict`/_`_gate_mask` treat
`net_crossings`/_`players_in_court`
227     as feature columns). **Seam 1 = injectable `WindowAnalyzer`s (forward)**: a registry
(mirrors `segmenter_registry`) of pure
228     per-window cues over the already-produced per-frame-aligned signals (WASB trajectory +
`PersonDetector.detections_per_frame`
229     + `rally_gate` net); each emits an `AnalyzerResult` = confidence-scored features ?
`candidate_segments.stats`, events ?
230     `events` table (buffer-on-candidate, flush-on-confirm), optional boundary hints; worked
example = a **service-fault**
231     analyzer (shuttle halts in net region ? `service_fault` signal + `net_hit` event + END
hint); injects at the FusionSegmenter
232     `_enrich_candidate_stats` hook, flag-gated default OFF. **Seam 2 = reconciliation pass
(backward)**: a report-first "linter
233     for rally decisions" ? `ConsistencyRule`s (trailing-dead-time?truncate *(owner's misfire
fix)*, no-net-crossing?drop,
234     over-merged?split, boundary-slack?tighten) flag decision-vs-evidence conflicts;
**report-first/apply-optional, NEVER silent**;
235     expressible as a harness `Variant` transform so LOVO-measurable on the 6 golden videos
BEFORE auto-apply; idempotent, fixed
236     precedence. ? **DESIGNED AGAINST the IN-FLIGHT fusion-P2 WIP** (uncommitted in the
shared tree: `fusion_hybrid.py`,

```

```

237     `fusion_features.py`, `FusionConfig.cue_flags`, `person_detector.detections_per_frame`,
`docs/QUALITY_ITERATIONS.md`) ? that
238     WIP was kept **byte-intact** (I isolated my single README line via git plumbing; tree is
CRLF / autocrlf=true, restored
239     owner README to exact session-start bytes 3841B). **Merge-order:** the doc cross-links
QUALITY_ITERATIONS.md + the fusion
240     hook, neither on master yet ? **merge #116 after (or with) fusion-P2** (noted in the PR
body). ?? NEXT: owner reviews #116;
241     implementation is the phased A1?A3 / R1?R3 plan in the doc (each ONE PR off master,
default OFF, promoted only on measured
242     LOVO lift) ? all owner-gated, none started. Ties [[candidate-segments-finetuning]] + the
harness ([[working-agreement-merging]]).
243
244     **Checkpoint:** 2026-06-13 (RALLY-QUALITY TRACK: HARNESS + PERSON-CUE EXTRACTION SHIPPED
? **PR #114 + #115 MERGED**,
245     master `df7d5df`; suite 421?441 green; 0 open PRs). **PR #115 (fusion P1):** extracted
the torchvision person detector
246     + ByteTrack + velocity windowing out of `yolo_hybrid` into
`backend/pipeline/detectors/person_detector.py`
247     (`PersonDetector`, a reusable cue producer); `yolo_hybrid` now ORCHESTRATES only
(chunk/prefetch/AI-handoff/DB) and
248     delegates Stage-1 to `self.person`. **No behavior change** ? proven by the unchanged
`process_video` output tests (mock
249     seam moved to `self.person`); pure detection unit tests ? new
`tests/test_person_detector.py`. person_detector is
250     mypy-CLEAN (`self.model: Any`) so it needed NO quarantine; the torchvision code that put
yolo_hybrid in quarantine left
251     it, but yolo_hybrid STAYS quarantined for PRE-EXISTING orchestration debt (implicit-
Optional process_video defaults,
252     Success|Failure ABC-override mismatch, a Future[...] annotation ? flagged in mypy.ini to
ratchet separately). ruff+mypy
253     clean, CI 4/4. Docs synced (fusion-plan P1?DONE, CODE_MAP person_detector entry +
yolo_hybrid-as-orchestration, NEXT_STEPS
254     item 2 ticked). ?? **NEXT (owner-steered):** **PR 3 = fusion P2.** The cheap, highest-
confidence, PURE-LOGIC piece =
255     **wire Gate A (`min_crossings`) into the live trajectory path**
(`trajectory_hybrid`/`wasb_hybrid`) + add the knob to
256     `IndexingConfig` ? **fixes the owner's held-shuttle false positive*, default OFF,
immediately A/B-able via the harness on
257     the 6 golden trajectories. Gate B (court-mask + player-count) is **owner-gated** on the
court-mask-source decision
258     (MULTI_SIGNAL_FUSION_PLAN §9.1: manual polygon vs auto homography ? owner). **P3 learned
fusion + the SxS-on-new-videos
259     RUN need GPU/owner-side** (persist person features on real video ? train on the 6 golden
? `compare_unlabeled(CONTROL,
260     +person)` on new footage). Two PRs this session within [[working-agreement-merging]]
self-merge authority.
261
262     **Checkpoint:** 2026-06-13 (EXPERIMENT HARNESS P1 SHIPPED ? **PR #114 MERGED**, master
`ecee786`; suite 421?440 green).
263     Picked up the rally-quality track (the desktop-agent pick-up `NEXT_STEPS` flagged; both
design docs `EXPERIMENT_HARNESS.md`
264     + `MULTI_SIGNAL_FUSION_PLAN.md` arrived via the session-start `git pull`
15e4e0e?9231af2, PR #112). Built
265     `backend/eval/experiment.py` = the thin variant/experiment layer from
`EXPERIMENT_HARNESS.md` P1 (the audit's "~80%
266     already exists" HELD ? verified file:line; pure glue over
metrics/calibration/classifier/training/regression.gt_hash/
267     gemini_refine.merge_overlaps/query_candidate_segments, **NO new scoring math**):
`Variant` (JSON recipe:
268     segmenter+config_overrides+cue_flags+Gate-A `min_crossings`+Gate-B `min_players`+learned
classifier; `spec_hash`; pinned
269     `CONTROL`), `compare` (labeled A/B, **LOVO-honest** train-on-others for learned
variants, B?A deltas + `label_set_hash`),
270     **`compare_unlabeled`** (the owner's **SxS-on-NEW-videos** ask ? agreed/a_only/b_only

```

```

via match_intervals, no GT),
271   `record_experiment` (? `eval_baselines/experiment_leaderboard.json`),
`load_video_candidates` (DB bridge). 19 tests incl.
272   the load-bearing **no-regression invariant** (all cues OFF ? byte-identical to CONTROL);
ruff+mypy clean; CI 4/4. Docs
273   synced (EXPERIMENT_HARNESS P1?DONE, CODE_MAP entry + stale regression-baseline-path fix,
NEXT_STEPS item 3 ticked). Owner
274   answered ramp-up Qs: **corpus = all 6** golden videos (he said "5";
LargeTestVideo/TestLargeVideo may be cuts of one match
275   ? kept all 6 per the manifest); **merge authority = self-merge on green then continue to
PR 2** (within [[working-agreement-merging]]).
276   ?? NEXT (this session): **PR 2 = fusion P1** person/vision-cue extraction (lift
torchvision+ByteTrack+velocity out of
277   `yolo_hybrid` ? `backend/pipeline/detectors/person_detector.py` producer; yolo_hybrid
consumes it, no behavior change ?
278   regression-safe). Then P3 learned fusion trains on the 6 golden videos +
`compare_unlabeled(CONTROL,+person)` SxS on new
279   footage (GPU/owner-side). Tracks [[candidate-segments-finetuning]] + [[collector-
indexer-bridge]].
280
281   **Checkpoint:** 2026-06-13 (WORKING-TREE CLEANUP ? master `15e4e0e`, tree now CLEAN).
Cleaned up the leftover
282   uncommitted bits (the cache-hygiene CODE was already MERGED as #109/#110 per the
checkpoint below ? nothing to rescue
283   there): (1) `config.json` GPU-run knobs **reverted to committed defaults** (was
wasb_hybrid/7200/8 ? gemini/1800/1;
284   values stay recorded in the cache-hygiene checkpoint below ? re-apply
`default_segementer: wasb_hybrid` in one line if
285   that's the intended new default; user didn't answer keep-vs-revert, took the safe
recoverable path). (2) stray root
286   `arial.ttf` (unused ? shipped watermark font is committed
`backend/assets/fonts/RobotoSlab-Regular.ttf` from #103) ?
287   moved to `scratch/`. (3) genuine untracked draft `docs/STORAGE_SHARING_MODEL.md`
(190-line storage-anchored
288   monetization model, owner co-design pending) ? branched + **PR #111 OPEN (owner-gated,
docs-only ? do NOT auto-merge**;
289   supersedes COMMERCIALIZATION.md framing if ratified ? ADR then). Master tree now shows
ONLY intentional carryovers
290   (untracked `scratch/` handoffs+repro scripts, `artifacts/` Gen-0 bundle) ? could be
gitignored for a pristine
291   `git status` if desired (offered, not done). Reminder: `python -m tools.cleanup_caches
--apply` reclaims the 96.5 GB
292   the SessionStart hook flagged.
293
294   **Checkpoint:** 2026-06-13 (CACHE HYGIENE shipped ? **PR #109 + docs #110 MERGED**,
master `15e4e0e`; + the ?5 reel +
295   manual cleanup). **#110** = diligence docs follow-up (README $Cache Hygiene + knobs +
tools/ line; CODE_MAP detector
296   self-clean, tools/cleanup_caches.py entry, common-task rows; min_shot_count API-parity
note) ? docs-only, CI 4/4 green.
297   Both repos now have **0 open PRs**. After the GX010125 E2E run left ~163 GB in WSL
`~/clips`, built a full cache-hygiene system (all on master):
298   (1) **self-cleaning runners** ? `indexing.{wasb,tracknet}.keep_frames` (default false)
deletes the frame cache +
299   staged video after a SUCCESSFUL run; on failure keeps them for resume; (2)
**tools/cleanup_caches.py** dry-run-first
300   ops tool (`--apply`/`--keep-video`/`--older-than`/`--include-wasbcache`/`--summary`);
(3) **disk guard** `min_free_gb`;
301   (4) **SessionStart nudge hook** in `.claude/settings.local.json` (local, gitignored).
Suite 409?421 green, ruff+mypy
302   clean. Policy saved ? [[cache-hygiene-policy]]. ? WSL-scan footgun: the owner's login
shell mangles `for e in `` loops
303   (loop var empty) ? use du/find arg-globbing. Also this session (earlier): merged
**#107** (API min_shot_count) + **#108**
304   (watermark Windows colon fix); ran GX010125 E2E on GPU ? **GX010125_highlights.mp4** (11

```

```

rallies >7 shots) +
305     **GX010125_highlights_min5.mp4** (24 rallies ?5, watermark now works ? validates #108)
via a stitch-only reuse script
306     (no re-inference); confirmed **Gemini got 1080p** (4K only at final stitch); delivered
an execution flowchart. Manual
307     WSL cleanup freed ~68 GB (kept GX010125's 93 GB frames + proxy per owner choice). ?
LOCAL uncommitted `config.json`:
308     default_segmenter=wasb_hybrid / min_shot_count=8 / wasb.timeout_sec=7200 from the run.
See [[gotcha-long-runs-gpu-wsl]].
309
310     **Checkpoint:** 2026-06-13 (SYNC + HANDOFF REFRESH ? read-only reconciliation by a ramp-
up session; no new code from
311     me this pass). `git pull --ff-only` = already up to date at **master `6de70aa` (#108)**.
Reconciled remote vs my last
312     session (#100): seven PRs landed from OTHER slates (NOT my work ? facts from git/gh):
**#101** configurable branding
313     watermark . **#102** thread config stitching paddings into live VideoCompiler sites .
**#103** bundle Apache-2.0
314     default watermark font + fallback log . **#104** C13 field kit (flyer/playbook/answer-
sheet/signals) . **#105** i18n
315     design plan (Kannada-first) . **#107** enforce `stitching.min_shot_count` on
/api/compile . **#108** escape Windows
316     drive-letter colon in ffmpeg filtergraph (fixes #103 watermark on Windows). **OPEN:
#106** docs/video-locality-model
317     (awaiting owner). ? **UNCOMMITTED WIP IN TREE (sibling/owner ? DO NOT CLOBBER):** a WSL
**cache-hygiene** slice in
318     progress, built on my `WslCommandMixin` ? `base.py` (`_wsl_free_gb`/`_wsl_rm_rf`),
`wasb_runner.py`+`tracknet_runner.py`
319     (`keep_frames`/`min_free_gb` config + pre-run disk guard + post-success cache delete),
`models.py` (new fields),
320     `config.json` (the GPU-run knobs: default_segmenter=wasb_hybrid, min_shot_count=8,
wasb.timeout_sec=7200), untracked
321     `arial.ttf` (watermark font) + `docs/STORAGE_SHARING_MODEL.md`. Not yet PR'd by whoever
owns it. [[session-handoff]]
322     rewritten to this state.
323
324     **Checkpoint:** 2026-06-13 (FIRST FULL E2E GPU RUN on a real 4K GoPro match + 2 merged
fixes).
325     Owner asked to run the indexer E2E on the GPU over `F:\GoPro Hero Black
12\KhelSutra\Badminton 12
326     June 2026\GX010125.MP4` (12.8-min 4K60, 11.5GB), highlights for rallies with >7 shots,
config-driven.
327     - **DELIVERED:** `output\GX010125_highlights.mp4` ? **11 rallies (shot_count 8?16) of 52
detected**;
328     true 4K 3840x2160 h264 59.94fps, 2:23, 1.15GB; stitched from the 4K ORIGINAL (proxy
only for indexing).
329     wasb_hybrid: 64.8% shuttle visibility, 18 candidate windows ? 52 Gemini-confirmed
rallies. min_shot_count=8
330     correctly excluded a 7-shot rally (>7 semantics). Gemini ran clean (no 503s);
~47-min wall clock.
331     - **2 PRs merged on green CI** (owner's [[working-agreement-merging]] standing
authorization):
332     **#107** `/api/compile` now enforces `stitching.min_shot_count` (parity w/
run_process_and_stitch);
333     **#108** sticher `_ff_quote` escapes the Windows drive-letter colon ? fixes the on-
by-default
334     bundled-font watermark (PR #103) silently failing on Windows (`fontfile='C:/...'`
broke ffmpeg
335     drawtext parsing; also fixed absolute logo `movie=` paths). Both with regression
tests; suite 409 green.
336     - **The hard-won run recipe is in [[gotcha-long-runs-gpu-wsl]]:** Modern Standby killed
the first 2
337     overnight attempts (keep-awake fix); cv2 PNG extraction (~38min) is the real
bottleneck not the GPU
338     (ffmpeg pre-extract, count-exact for CFR ? skips it); 1080p-proxy-index / 4K-stitch;

```

```

per-phase
339     observability (cache manifest for WASB, handoff dir for Gemini). Reusable scripts in
`scratch/`
340     (run_gx010125.py, keep_awake.ps1, probe_*.py).
341     - ? LOCAL UNCOMMITTED `config.json`: default_segmenter=wasb_hybrid, min_shot_count=8,
wasb.timeout_sec=7200
342     (set for this run; revert to gemini/1/1800 if not wanted). Watermark fix NOT applied
to the delivered
343     reel (it ran before #108); a re-stitch would hit the DB segment cache + now-working
watermark (~5min).
344
345     **Checkpoint:** 2026-06-12 (latest ? PAID COMPONENT REMOVED + VIDEO-LOCALITY MODEL; **PR
#104 updated `9c84fea` +
346     PR #106 OPEN `85abe26`, both CI green, both AWAITING OWNER**). Owner spec: (1) remove
ALL associate-pay from the
347     field kit ("something else in mind") ? flyer/playbook now "Engagement terms: discussed
individually"; honesty
348     mechanics reworded ("judged on honest capture, not positive answers"); KEPT: WTP
research questions (study subject)
349     + the ACADEMY's Phase-2 session fee (different party ? flagged, owner may strip too).
(2) Storage/upload questions
350     answered from code: uploads ARE stored in full (`security.upload_dir`, NO retention
policy ? BACKLOG); 10GB CANNOT
351     upload (max_upload_mb=4096 rejects at init; ?90MB sequential parts; hours on Indian
uplinks). (3) NEW
352     `docs/VIDEO_LOCALITY_MODEL.md` (#106): central insight = timestamps are the product,
video is dead weight (only the
353     compiler touches original bytes, cuttable where the original lives); locality ladder
L0?L4 (L2 proxy-upload reuses
354     the collector ADR-002 frame-identical contract; L1-direct = client?Gemini, zero bytes
through us; L0 = owned-model
355     north star; pilot = L4 sneakernet); **frugality frame (owner): bills affordable iff GPU
doesn't burn ? L2/L1 serving
356     is GPU-less by construction, Gemini €/hr, GPU stays on the owned 2070S, per-video cost
ceiling required**
357     **academy-visit dividend: Phase-2 consented recordings ARE the multi-court target
corpus, show-back = live L2 dress
358     rehearsal, demo reactions = build backlog**
359     10 open questions OQ1?OQ10 (retention, MD5
identity, Gemini consent,
359     key custody, packaging, review source, reel delivery, uplink data, volume, L3 survival);
rec lean = L2 default
360     (NOT a greenlight). NB: #106's base commit was pushed by owner/sibling mid-session
(shared checkout). ?? owner:
361     review/merge #104 + #106; decide academy-fee keep/strip; the "something else" comp
structure is his to define.
362
363     **Checkpoint:** 2026-06-12 (latest ? C13 KIT REFRAMED PRODUCT-FEEDBACK-FIRST per owner
spec + **RECRUITING FLYER
364     added; PR #104 UPDATED `985a329`, CI 4/4 green, AWAITING OWNER REVIEW** ? do NOT merge
without him). Owner spec:
365     hire someone to visit academies, talk to coaches AND students (how does it help +
improvement feedback); stretch =
366     record ? pipeline ? show the academy ITS OWN sample result; UI/hosting "ready by then".
Kit now = 4 docs by audience:
367     **`PRODUCT_FIELD_ASSOCIATE_FLYER.md`** (NEW ? candidate-facing recruiting flyer, owner
fills ?/phone + recruits) .
368     **`FIELD_VISIT_PLAYBOOK.md`** (RENAMED from FIELD_RESEARCH_ASSOCIATE_FLYER ? day-one
playbook: demo-after-Q6 rule,
369     adult-18+-students-only w/ coach-points-out gating, Phase-2 show-back + scripted "can we
keep using it?" no-promise
370     reply) . answer sheet (+demo-reaction/student blocks) . PMF_FIELD_SIGNALS (verdict
machinery UNCHANGED; demo/student
371     data = product feedback only). 2nd adversarial critique (`w5aexakz2`, 30 agents): 12
confirmed+fixed (key honesty
372     fixes: "web app IS opening"? "we're building?will be able to try"; "match report"? "simple

```


summary"; "finds EVERY
373 rally" dropped; minors loophole in "senior/adult students" closed), 15 refuted. ??
owner: review #104 (thresholds
374 \$3 + flyer copy) ? merge ? fill ?/phone ? trim 60s demo reel ? recruit.
375 Owner ratified in-chat: **PMF validation = physically visiting Bangalore academies and
reading signals.** Built the
376 decision layer on top of the existing `docs/FIELD\_RESEARCH\_ASSOCIATE\_FLYER.md` (which IS
the "recruit hand-out" the
377 owner remembered; the khelacharya PDFs in OneDrive are the separate internship flyer ?
see
378 [[sibling-internship-repo]]: **`docs/PMF\_FIELD\_SIGNALS.md`** (H1-H7 ? script-question
map, G/A/R anchors, H5
379 counting rule = pre-ladder-or-budget-converted ??300 only ? anchored ?300 NEVER counts,
pinned stopping rule n=15
380 target/logistics-only/verdict-once, numeric verdicts PROCEED / PIVOT A record+index /
PIVOT B partner-API
381 (PB Vision*PodPlay footnote now in-repo) / PIVOT C parent-payer / FALSIFIED) +
`docs/FIELD\_VISIT\_ANSWER\_SHEET.md`
382 (printable per-visit form: verbatim answers, required most-negative quote, before/after-
ladder tick, OFFICE-USE
383 score grid) + flyer Q5 rewritten open-ended-first (old wording read the ?100/300/500
ladder ALOUD ? center-bias
384 manufactures the exact ?300 PROCEED threshold; caught by the 38-agent adversarial
critique `ww0rvaqf3`: 1 high/2 med/
385 2 low confirmed+fixed, 9 refuted) + Stupa added to listen-for + NEXT_STEPS item 5 = kit
COMPLETE. ?? owner: review
386 #104 thresholds ? merge ? fill flyer ?/phone placeholders ? recruit the associate ?
first 2 visits paired. PDFs of
387 flyer+answer-sheet: generate AFTER placeholders are filled (offer stands).
388
389 **Checkpoint:** 2026-06-12 (latest ? ROORA VENDOR LOOP COMPLETE: **collector #23+#24+#25
MERGED** sequentially on
390 green CI, owner pre-authorized; master `ca61405`, 0 open PRs, only owner's `.gitignore`
edit uncommitted).
391 **#23 boundary_grid_snap BLOCKER** in detect_issues (gcd-of-boundary-frame-diffs > ±0.3s
bar ? gate BLOCKED; skips
392 CSV-path/no-fps/<3 rallies) ? ? **replayed against the REAL 6-June PB+TT pilot zips:
both auto-block with "All 12
393 rally boundaries sit on a 30-frame grid (~0.50s @ 59.94fps)"** (the manual forensics,
now automated) + ROORA_BRIEF
394 Phase-B revision (\$7 REQUIRED method: CVAT step 1, first/last box on exact contact/dead
frame, sparse interior OK;
395 \$4 grid-snap auto-detected warning; \$2 skip-rallies-in-progress-at-file-edges; \$1
annotation-optimized-copy note).
396 **#24 prep-annotation-batch**: mirrors `<input-dir>/<match-folder>/<video>` ? proxies-
only tree w/ same layout
397 (thin orchestrator over the hardened single flow; per-video isolation; resumable "have"
skips; live 2-court E2E
398 smoke green). **#25 ROORA_RUNBOOK.md** (docs/annotation/): operator E2E ? send checklist
(layout?batch?manifest
399 commit?intake+email w/ the 2 non-negotiables), receive (gate decision table incl. grid-
snap, review sheet,
400 import/export on ORIGINALS), troubleshooting table for every guard. Suite **180 green**.
?? Vendor loop is now fully
401 toolled+documented. ?? AGREED NEXT (owner, session wrap 2026-06-12): (1) owner fills the
canonical doc's TODO(avidullu)
402 markers + screenshots, labels it **v8**;

(2) **joint TEST RUN of the full vendor loop on
real footage BEFORE
403 anything goes to Roora** (runbook \$1 send-side end-to-end: batch layout ? prep-
annotation-batch ? manifest ?
404 mock-receive via validate-delivery on an old delivery); (3) only then send. **v7 ANOMALY
RESOLVED (2026-06-12):**
405 the canonical vendor-facing Brief+Guide is the Google Doc at
docs.google.com/document/d/1NvAXE0q7t4Z0QddLkEfnta0DU982yDZaqkGNPMe8EY0
406 (provenance: generated from master `ca61405`; supersedes Drive v3/v5/v6/v7; owner

confirmed superset-of-v7) ? see
407 [[roora-guidelines-drive-doc]]; before sending: pick ONE version label + fill its
TODO(avidullu) markers/screenshots.
408
409 **Checkpoint:** 2026-06-12 (earlier ? ADR-002 IMPLEMENTED: **collector #22 MERGED**
`a068fd1`, owner authorized
410 merge-on-green; #21 ADR + #20 + indexer #102/#103 all merged earlier today, slate was
clean). Shipped: `curate
411 prep-annotation-proxy` + `src/core/proxy.py` ? ?1080p/-g 15`/CRF 19/audio-kept re-
encode w/ `-fps\_mode passthrough`
412 (NEVER `-r`), post-encode parity verify (exact frame count; fps $\pm 0.1\%$; duration like-
with-like stream/container),
413 delete-on-ANY-failure (incl. Ctrl-C), VFR refusal w/o `--allow-vfr`, provenance ?
414 `data/annotation\_proxy\_manifest.csv` (source+proxy MD5s; video_id collision guard BEFORE
encode; Excel-lock prints
415 row + keeps proxy), same-file guard normcase/realpath/samefile (NTFS case-variant
`--out`+`--force` would have let
416 `ffmpeg -y` TRUNCATE the source ? review verifier reproduced it live pre-fix). Suite
169 green (+32), CI green.
417 Verified live on ffmpeg 8.1.1 (smoke clip 59.94fps: parity 179 frames both files; smoke
caught `-vsync` deprecation ?
418 `-fps\_mode`). Process: ultracode adversarial review workflow (`wyv8524zf`, 18 agents)
confirmed 1 high/2 med/7 low ?
419 all fixed pre-merge. Docs: INGESTION_PIPELINE step 0 + diagram leg + onboarding FAQ +
sampling-FAQ split
420 (highlights-vs-golden), ADR-002 provenance bullet as-implemented amendment,
PIPELINE_ROLE ADR refs qualified.
421 ?? REMAINING ADR-002 follow-ups (separate items, NOT started): Phase-B brief revision
(frame step 1 / exact boundary
422 frames / "we supply proxies") ? owner words vendor comms; grid-snap QA check in
validate-delivery. First real use:
423 run prep-annotation-proxy on the next Roora batch before handoff.
424
425 **Checkpoint:** 2026-06-12 (earlier ? ANNOTATION-PROXY DECISION RATIFIED ? **collector
PR #21 OPEN**, docs-only
426 **ADR-002** in collector `docs/DECISIONS.md`). Owner agreed in-chat w/ my rec from the
assessment checkpoint below:
427 proxies generated in-repo via a future `curate prep-annotation-proxy`. ADR records the
proxy contract (SAME fps ?
428 never reduce; no trim; frame-count parity check; 1080p/CRF18-20/-g15/audio kept; intake-
manifest row w/ original+proxy
429 MD5s), the why-collector rationale (owns the vendor flow = command #0 of convert-
cvat?validate-delivery?import-local;
430 verification+provenance is the durable 80%, one-off scripts skip it), and the **Phase-B
necessity argument the owner
431 wanted captured: ~12-15 videos/sport $\times$ ~30 min 4K/59.94 = >100k frames/file ? vendor
frame-step-1 annotation
432 infeasible without proxies ? near-precondition, not optimization**. Follow-ups recorded
IN the ADR: Phase-B brief
433 convention line (step=1, first/last box on exact contact/dead frame) + re-add grid-snap
QA to validate-delivery.
434 Collector checkout back on master (owner's stray `.gitignore` edit untouched). ? sibling
PR #20 (below) edits
435 rally_cvat.py docstring rationale 0.5/1.0?0.5/0.5 ? no file overlap w/ #21, both docs-
true. ?? owner reviews/merges
436 #21 (+ #102/#20 below); subcommand IMPLEMENTATION stays owner-gated (separate PR on
greenlight).
437
438 **Checkpoint:** 2026-06-12 (stitching-padding wiring fix ? **indexer #102 + collector
#20 OPEN, CI green, awaiting
439 owner review**). Audit finding (d) (2026-06-10) fixed: config.json
`stitching.{begin,end}\_padding` was silently
440 ignored by the live compiler ? both `backend/main.py` VideoCompiler sites
(`/api/compile` + CLI trim) passed only
441 `watermark=`, so reels always used ctor defaults 0.5/**1.0** vs configured 0.5/0.5. Fix

```

threads the typed
442     StitchingConfig paddings into both sites (+ `StitchingConfig` re-export from
backend.config); regression test drives
443     the REAL compiler through /api/compile (mocked ffmpeg) asserting a config override
(2.0/3.0) reaches `-ss/-to`
444     (8.0/17.0 for a 10?14s segment ? unwired defaults would give 9.5/15.0); CODE_MAP truth-
syncd (ctor-vs-config
445     distinction + recipe row now cites the main.py wiring). ? **Behavior change on merge:
live API reels end padding
446     1.0?0.5s** (the config value). Companion **collector #20** = docstring-only:
`rally_cvat.py::detect_issues` cited the
447     hardcoded 0.5/1.0 + a stale module path; now cites config-driven 0.5/0.5 (rationale
still holds; ? cross-checks the
448     sibling Roora-proxy thread below ? boundary-slack rationale now assumes 0.5s end
padding, not 1.0s). Verified: indexer
449     suite **401 green** + 4/4 CI lanes; collector **137 green** + CI. Both repos back on
master; owner's uncommitted
450     collector `.gitignore` edit untouched. NOTE: untracked `arial.ttf` at indexer root (not
mine ? left alone, likely a
451     watermark-font experiment). ?? owner reviews/merges #102 + collector #20.
452

```

7. assistant (2026-06-14T09:49:43.199Z)

Let me ramp up on the eval harness internals and a sample golden JSON. I'll read the key files in parallel.

8. user (2026-06-14T09:49:44.170Z)

```

1     """Pure-CPU LOVO comparison: shuttle-only vs shuttle+person learned fusion on the golden
set.
2
3     The fusion-P3 measurement. Loads the per-video feature JSONs written by
`fusion_golden.py`
4     and runs the SHIPPED experiment harness (`experiment.compare`, honest leave-one-video-
out) ?
5     no new scoring/labeling math. The two variants differ ONLY by `feature_names` (the
harness's
6     designed seam): the person features are extra columns the learned LogisticRegression
weights,
7     so the held-out ?F1 is an honest answer to "does the person cue help?" ? not a hunch.
8
9     Run (no GPU):
10     python -m backend.eval.fusion_compare --features-dir output --record
11     """
12     from __future__ import annotations
13
14     import argparse
15     import glob
16     import json
17     import os
18     from typing import Any, Dict, List
19
20     from backend.eval import experiment
21     from backend.eval.eval_stats import paired_delta_ci
22     from backend.eval.fusion_golden import PERSON_FEATURE_NAMES, SHUTTLE_FEATURE_NAMES
23
24
25     def delta_stats(res: Dict[str, Any], metric: str = "f1") -> Dict[str, Any]:
26         """Paired (by video) B?A summary for one metric: mean delta + bootstrap CI + sign
test
27         (eval_stats, C1 ? "never a bare mean"). Aligns the two variants' per-video held-out
scores
28         by video name, so the headline ?F1 comes with how-wide + did-it-win-on-enough-
videos."""

```

```

29     by_a = {p["video"]: p[metric] for p in res["variant_a"]["per_video"]}
30     by_b = {p["video"]: p[metric] for p in res["variant_b"]["per_video"]}
31     vids = [p["video"] for p in res["variant_a"]["per_video"]]
32     return paired_delta_ci([by_a[v] for v in vids], [by_b[v] for v in vids])
33
34
35 def load_videos(features_dir: str) -> List[experiment.VideoCandidates]:
36     """Load every ``*_golden_features.json`` into a VideoCandidates (intervals + the 10
37     feature columns + golden gts). The harness derives labels from gts itself."""
38     videos: List[experiment.VideoCandidates] = []
39     for path in sorted(glob.glob(os.path.join(features_dir, "*_golden_features.json"))):
40         with open(path, encoding="utf-8") as f:
41             d = json.load(f)
42             cand = d["candidates"]
43             intervals = [(float(c["start"]), float(c["end"])) for c in cand]
44             features: Dict[str, List[float]] = {}
45             for fn in SHUTTLE_FEATURE_NAMES:
46                 features[fn] = [float(c["shuttle"][fn]) for c in cand]
47             for fn in PERSON_FEATURE_NAMES:
48                 features[fn] = [float(c["person"][fn]) for c in cand]
49             gts = [(float(a), float(b)) for a, b in d["gts"]]
50             videos.append(experiment.VideoCandidates(name=d["name"], intervals=intervals,
51                                                         features=features, gts=gts))
52     return videos
53
54
55 def compare(videos: List[experiment.VideoCandidates], iou: float = 0.5,
56             threshold: float = 0.5) -> Dict[str, Any]:
57     shuttle_only = experiment.Variant(name="shuttle_only",
58                                       feature_names=list(SHUTTLE_FEATURE_NAMES),
59                                       threshold=threshold)
60     fusion = experiment.Variant(name="shuttle_person_fusion",
61                                 feature_names=list(SHUTTLE_FEATURE_NAMES) +
62                                 list(PERSON_FEATURE_NAMES),
63                                 threshold=threshold)
64     return experiment.compare(videos, shuttle_only, fusion, iou=iou, honest=True)
65
66
67 def format_result(res: Dict[str, Any], iou: float) -> str:
68     a, b, d = res["variant_a"]["aggregate"], res["variant_b"]["aggregate"], res["delta"]
69     ds = delta_stats(res)
70     st = ds["sign_test"]
71     lines = [
72         f"=== Fusion P3 LOVO (n={res['num_videos']} golden videos, IoU={iou}, honest) ===",
73         f" shuttle-only : F1={a['f1']:.3f} P={a['precision']:.3f} R={a['recall']:.3f} IoU={a['mean_iou']:.3f}",
74         f" +person : F1={b['f1']:.3f} P={b['precision']:.3f} R={b['recall']:.3f} IoU={b['mean_iou']:.3f}",
75         f" ?F1 = {d['f1']:.3f} ({res['delta_pct']['f1']:.1f}%)",
76         f"?P={d['precision']:.3f} ?R={d['recall']:.3f}",
77         f" paired ?F1 {ds['mean_delta']:.3f} [95% CI {ds['ci_low']:.3f}, {ds['ci_high']:.3f}] ",
78         f"| sign test {int(st['wins'])}W/{int(st['losses'])}L (p={st['p_value']:.3f}) ",
79         f"| CI clears 0: {ds['ci_excludes_zero']}",
80         f" per-video ?F1:",
81     ]
82     for pv in res["per_video_delta"]:
83         lines.append(f" {pv['video']}<26s ?F1={pv['f1']:.3f}")
84     return "\n".join(lines)
85
86
87 def main() -> None:
88     ap = argparse.ArgumentParser(description="LOVO fusion lift: shuttle-only vs shuttle+person (golden set).")

```

```

86         ap.add_argument("--features-dir", default="output")
87         ap.add_argument("--iou", type=float, default=0.5)
88         ap.add_argument("--record", action="store_true", help="append the result to the
experiment leaderboard")
89         args = ap.parse_args()
90
91         videos = load_videos(args.features_dir)
92         if len(videos) < 2:
93             raise SystemExit(f"need >=2 *_golden_features.json in {args.features_dir} (found
{len(videos)})")
94         res = compare(videos, iou=args.iou)
95         print(format_result(res, args.iou))
96         if args.record:
97             row = experiment.record_experiment(res)
98             print(f"recorded -> eval_baselines/experiment_leaderboard.json
({row['timestamp']}")
99
100
101     if __name__ == "__main__":
102         main()
103

```

9. user (2026-06-14T09:49:44.659Z)

```

1     """Fusion P4 ? measure the INCREMENTAL audio-cue lift on top of the person cue (CPU-
only).
2
3     Reuses the person-feature JSONs from `fusion_golden.py` (no GPU re-detect): `augment`
runs
4     `detect_hits_in_video` once per video (ffmpeg + numpy, CPU) and adds the 3 audio
features per
5     window; `compare_increment` then runs the SHIPPED harness between a shuttle+person
variant
6     and a shuttle+person+audio variant ? so audio is measured as a marginal contribution ON
TOP
7     of person, not against raw shuttle. Audio enters only as columns the learned scorer
weights
8     (no special-casing, default-OFF); the golden-set ?F1 decides whether it earns its keep.
9
10    Run (no GPU):
11        python -m backend.eval.fusion_audio --manifest output/human_lovo_manifest.json
--features-dir output --augment --compare
12    """
13    from __future__ import annotations
14
15    import argparse
16    import glob
17    import json
18    import os
19    from typing import Any, Dict, List
20
21    from backend.eval import experiment
22    from backend.eval.fusion_compare import delta_stats
23    from backend.eval.fusion_golden import PERSON_FEATURE_NAMES, SHUTTLE_FEATURE_NAMES,
_derive_video_path
24    from backend.pipeline.detectors import audio_features as af
25    from backend.pipeline.audio.hit_detector import detect_hits_in_video
26
27
28    def augment(manifest_path: str, features_dir: str, k: float = 2.5) -> List[str]:
29        """Add the 3 audio features to each existing ``<name>_golden_features.json``
window."""
30        with open(manifest_path, encoding="utf-8") as f:
31            specs = json.load(f)
32            written: List[str] = []

```

```

33     for spec in specs:
34         jp = os.path.join(features_dir, f"{spec['name']}_golden_features.json")
35         if not os.path.isfile(jp):
36             print(f"[fusion_audio] {spec['name']}: no feature JSON ? run fusion_golden
first; skipping")
37             continue
38         video = _derive_video_path(spec)
39         print(f"[fusion_audio] {spec['name']}: detecting audio hits (k={k})...",
flush=True)
40         hits = detect_hits_in_video(video, k=k)
41         with open(jp, encoding="utf-8") as f:
42             data = json.load(f)
43             for c in data["candidates"]:
44                 c["audio"] = af.compute_audio_features(hits, c["start"], c["end"])
45             with open(jp, "w", encoding="utf-8") as f:
46                 json.dump(data, f, indent=2)
47             print(f"[fusion_audio] {spec['name']}: {len(hits)} hits over
{len(data['candidates'])} windows", flush=True)
48             written.append(jp)
49         return written
50
51
52 def load_videos_with_audio(features_dir: str) -> List[experiment.VideoCandidates]:
53     """Load the feature JSONs into VideoCandidates including the audio columns (requires
54     `augment` to have run ? raises if a JSON lacks audio)."""
55     videos: List[experiment.VideoCandidates] = []
56     for path in sorted(glob.glob(os.path.join(features_dir, "*_golden_features.json"))):
57         with open(path, encoding="utf-8") as f:
58             d = json.load(f)
59             cands = d["candidates"]
60             if cands and "audio" not in cands[0]:
61                 raise SystemExit(f"{path} has no audio features ? run `--augment` first")
62             intervals = [(float(c["start"]), float(c["end"])) for c in cands]
63             features: Dict[str, List[float]] = {}
64             for fn in SHUTTLE_FEATURE_NAMES:
65                 features[fn] = [float(c["shuttle"][fn]) for c in cands]
66             for fn in PERSON_FEATURE_NAMES:
67                 features[fn] = [float(c["person"][fn]) for c in cands]
68             for fn in af.AUDIO_FEATURE_NAMES:
69                 features[fn] = [float(c["audio"][fn]) for c in cands]
70             gts = [(float(a), float(b)) for a, b in d["gts"]]
71             videos.append(experiment.VideoCandidates(name=d["name"], intervals=intervals,
72                                                         features=features, gts=gts))
73     return videos
74
75
76 def compare_increment(videos: List[experiment.VideoCandidates], iou: float = 0.5,
77                       threshold: float = 0.5) -> Dict[str, Any]:
78     """Incremental audio lift: shuttle+person vs shuttle+person+audio (honest
LOVO)."""
79     person = experiment.Variant(
80         name="shuttle_person",
81         feature_names=list(SHUTTLE_FEATURE_NAMES) + list(PERSON_FEATURE_NAMES),
threshold=threshold)
82     person_audio = experiment.Variant(
83         name="shuttle_person_audio",
84         feature_names=list(SHUTTLE_FEATURE_NAMES) + list(PERSON_FEATURE_NAMES) +
list(af.AUDIO_FEATURE_NAMES),
85         threshold=threshold)
86     return experiment.compare(videos, person, person_audio, iou=iou, honest=True)
87
88
89 def format_result(res: Dict[str, Any], iou: float) -> str:
90     a, b, d = res["variant_a"]["aggregate"], res["variant_b"]["aggregate"], res["delta"]
91     ds = delta_stats(res)

```

```

92         st = ds["sign_test"]
93         lines = [
94             f"=== Fusion P4 incremental AUDIO LOVO (n={res['num_videos']}, IoU>={iou},
honest) ===",
95             f"  shuttle+person          : F1={a['f1']:.3f}  P={a['precision']:.3f}
R={a['recall']:.3f}",
96             f"  shuttle+person+audio : F1={b['f1']:.3f}  P={b['precision']:.3f}
R={b['recall']:.3f}",
97             f"  incremental ?F1 = {d['f1']:.3f}  ({res['delta_pct']['f1']:.1f}%)",
98             f"  paired ?F1 {ds['mean_delta']:.3f}  [95% CI {ds['ci_low']:.3f},
{ds['ci_high']:.3f}]  "
99             f"| sign test {int(st['wins'])}W/{int(st['losses'])}L (p={st['p_value']:.3f})  "
100            f"| CI clears 0: {ds['ci_excludes_zero']}",
101        ]
102        for pv in res["per_video_delta"]:
103            lines.append(f"      {pv['video']:<26s} ?F1={pv['f1']:.3f}")
104        return "\n".join(lines)
105
106
107    def main() -> None:
108        ap = argparse.ArgumentParser(description="Fusion P4 ? incremental audio lift on the
golden set.")
109        ap.add_argument("--manifest", default="output/human_lovo_manifest.json")
110        ap.add_argument("--features-dir", default="output")
111        ap.add_argument("--k", type=float, default=2.5, help="hit-detector sensitivity")
112        ap.add_argument("--augment", action="store_true", help="add audio features to the
existing JSONs")
113        ap.add_argument("--compare", action="store_true", help="report the incremental audio
LOVO ?F1")
114        ap.add_argument("--iou", type=float, default=0.5)
115        ap.add_argument("--record", action="store_true")
116        args = ap.parse_args()
117
118        if args.augment:
119            augment(args.manifest, args.features_dir, args.k)
120        if args.compare:
121            videos = load_videos_with_audio(args.features_dir)
122            if len(videos) < 2:
123                raise SystemExit(f"need >=2 augmented feature files in {args.features_dir}")
124            res = compare_increment(videos, iou=args.iou)
125            print(format_result(res, args.iou))
126            if args.record:
127                row = experiment.record_experiment(res)
128                print(f"recorded -> eval_baselines/experiment_leaderboard.json
({row['timestamp']}")
129
130
131    if __name__ == "__main__":
132        main()
133

```

10. user (2026-06-14T09:49:45.709Z)

```

1    """Experiment harness ? A/B + shadow + side-by-side for rally detection.
2
3    This is the thin **variant/experiment layer** that `docs/EXPERIMENT_HARNESS.md`
4    identifies as the one missing piece: ~80% of the machinery (scoring, LOVO, the
5    no-regression gate, the pinnable classifier artifact, the persisted-candidate
6    substrate) already exists ? this module composes it. **No new scoring math lives
7    here**; everything defers to:
8
9    - `backend.eval.metrics`      ? `score`, `match_intervals`, `interval_iou`
10   - `backend.eval.training`     ? `label_candidates` (y by the same IoU matcher)
11   - `backend.eval.classifier`   ? `LogisticRegression` (the pinnable JSON model)
12   - `backend.eval.gemini_refine` ? `merge_overlaps` (the blessed window merge)

```

```

13     - `backend.eval.regression`      ? `gt_hash` (label-set fingerprint)
14     - `backend.eval.calibration`     ? `pct_improvement`
15
16 The thesis (`EXPERIMENT_HARNESS.md` §2): separate the expensive, risky DETECTION
17 (per-frame signals ? run once on the GPU/WSL box, persisted into
18 `candidate_segments.stats`) from the cheap, swappable DECISION (given those
19 signals, where are the rallies). A `Variant` is a declarative re-scoring recipe;
20 given persisted candidate features it produces `List[Interval]` deterministically,
21 with no GPU and zero production risk. So most experiments are pure offline
22 re-scoring of stored data and never touch the served pipeline.
23
24 Three modes (`EXPERIMENT_HARNESS.md` §5):
25     - `compare()`      ? labeled A/B vs golden labels: per-video + LOVO-honest
26       aggregate deltas. Mirrors `distill_local`'s honest train-on-others/predict-
27       held-out loop, but variant-parameterised.
28     - `compare_unlabeled()` ? SxS on NEW footage with no ground truth: where do two
29       variants agree / diverge (A-only, B-only, agreed). The divergences are what a
30       human spot-checks (and what two highlight reels would show side by side).
31     - the promotion gate stays `run_regression.py` + `regression_baseline.json`
32       (exit 0/1/3); rollback = flip the variant/config, never `git revert`.
33
34 Pure + offline + unit-tested. The only DB-touching helper is
35 `load_video_candidates`, which reads persisted candidates via
36 `Database.query_candidate_segments`.
37 """
38 from __future__ import annotations
39
40 import json
41 import logging
42 import os
43 from dataclasses import dataclass, field
44 from datetime import datetime, timezone
45 from hashlib import sha1
46 from typing import Any, Callable, Dict, List, Optional, Sequence
47
48 from backend.eval.calibration import pct_improvement
49 from backend.eval.classifier import LogisticRegression
50 from backend.eval.gemini_refine import merge_overlaps
51 from backend.eval.metrics import Interval, match_intervals, score
52 from backend.eval.regression import gt_hash
53 from backend.eval.training import label_candidates
54 from backend.eval import eval_stats as _stats          # C1: CIs + paired sign test
55 from backend.eval import segmentation_metrics as _segm  # C4: F1@k + segment-count
56 ratio
57 scores
58
59 from backend.eval.score_calibration import fit_isotonic, fit_platt # C2: calibrate cue
60
61 logger = logging.getLogger(__name__)
62
63 # Where a comparison run appends one reproducible record. Tracked location (NOT under
64 # the gitignored output/) so the leaderboard survives a clean checkout, mirroring
65 # regression.DEFAULT_BASELINE_PATH.
66 DEFAULT_LEADERBOARD_PATH = os.path.join("eval_baselines", "experiment_leaderboard.json")
67 _METRICS = ("precision", "recall", "f1", "mean_iou")
68
69
70 class ExperimentError(ValueError):
71     """A variant cannot be evaluated as configured (e.g. a learned variant asked to
72     score an unlabeled video with no pinned model, or a gate referencing an absent
73     feature)."""
74
75 # ----- Variant

```



```

76 @dataclass
77 class Variant:
78     """A declarative re-scoring recipe ? NOT a code branch (`EXPERIMENT_HARNESS.md` §4).
79
80     A variant turns one video's persisted candidate windows + features into rally
81     intervals. Every knob defaults to the control behaviour (no gate, no learned
82     filter), so a variant that merely carries a new, disabled cue is byte-identical
83     to control ? the core no-regression guarantee.
84
85     Fields:
86     - ``segmenter`` / ``config_overrides`` / ``cue_flags`` ? informational provenance
87       for the leaderboard and for selecting the served path (the existing
88       ``segmenter_registry`` + ``IndexingConfig`` do the actual selection in production).
89       New cues are recorded here default-OFF.
90     - ``min_crossings`` ? decision-gate Gate A: keep only windows whose persisted
91       ``net_crossings`` feature is >= this. 0 = off. This is the eval-only gate from
92       ``rally_gate.py`` expressed as a re-scoring knob (kills service-feed / held-
shuttle
93       windows ? see ``MULTI_SIGNAL_FUSION_PLAN.md` §3.1).
94     - ``min_players`` ? decision-gate Gate B (the future person cue): keep windows
95       whose persisted ``players_in_court`` is >= this. 0 = off (no-op until the person
96       cue persists that feature).
97     - ``classifier_model`` ? path to a frozen ``LogisticRegression`` JSON. When set,
98       ``compare()`` scores videos with the SHIPPED model as-is (no per-fold training);
99       required for ``compare_unlabeled`` on a learned variant.
100    - ``feature_names`` ? the persisted features the learned filter consumes. When set
101      WITHOUT ``classifier_model``, ``compare()`` trains a fresh model per LOVO fold
102      (honest) ? the recipe, not a pinned artifact.
103    - ``threshold`` ? classifier probability cutoff. ``merge_gap`` ? post-filter
104      overlap-merge gap (seconds), the same ``gemini_refine.merge_overlaps`` default.
105    """
106
107    name: str
108    segmenter: str = "wasb_hybrid"
109    config_overrides: Dict[str, Any] = field(default_factory=dict)
110    cue_flags: Dict[str, bool] = field(default_factory=dict)
111    min_crossings: int = 0
112    min_players: int = 0
113    classifier_model: Optional[str] = None
114    feature_names: Optional[List[str]] = None
115    threshold: float = 0.5
116    merge_gap: float = 1.0
117    # C2 / threshold-in-LOVO (default OFF ? unchanged behaviour). When set, the
operating
118    # point is chosen on the TRAINING fold, never the held-out one ? fixing the
first-A/B
119    # failure where a fixed 0.5 zeroed out a small-candidate video.
120    tune_threshold: bool = False          # pick the F1-best threshold on the train fold
121    calibrate: Optional[str] = None        # None | "platt" | "isotonic" ? calibrate
proba first
122
123    def is_learned(self) -> bool:
124        """True if this variant applies a learned classifier (pinned model or
recipe)."""
125        return self.classifier_model is not None or bool(self.feature_names)
126
127    def to_dict(self) -> dict:
128        return {
129            "name": self.name,
130            "segmenter": self.segmenter,
131            "config_overrides": self.config_overrides,
132            "cue_flags": self.cue_flags,
133            "min_crossings": self.min_crossings,
134            "min_players": self.min_players,
135            "classifier_model": self.classifier_model,

```

```

136         "feature_names": self.feature_names,
137         "threshold": self.threshold,
138         "merge_gap": self.merge_gap,
139         "tune_threshold": self.tune_threshold,
140         "calibrate": self.calibrate,
141     }
142
143     @classmethod
144     def from_dict(cls, d: dict) -> "Variant":
145         known = {f for f in cls.__dataclass_fields__} # type: ignore[attr-defined]
146         return cls(**{k: v for k, v in d.items() if k in known})
147
148     def spec_hash(self) -> str:
149         """Stable 12-char hash of the recipe ? identifies the variant in the leaderboard
150         and makes a result reproducible. The pinned **control** variant always hashes
151         the
152         same, so a regression is always traceable to a recipe change, never to drift."""
153         blob = json.dumps(self.to_dict(), sort_keys=True, default=str)
154         return sha1(blob.encode()).hexdigest()[:12]
155
156 #: The pinned, frozen baseline: candidates as detected, merged, no gate, no learned
157 #: filter. Always the comparison reference; "rollback" = make this the served variant.
158 CONTROL = Variant(name="control")
159
160
161 # ----- persisted candidates
162
163
164 @dataclass
165 class VideoCandidates:
166     """One video's persisted detection, ready for offline re-scoring.
167
168     ``intervals`` are the candidate windows; ``features`` is column-major
169     (``feature_name -> per-candidate value``, each list aligned to ``intervals``);
170     ``gts`` are golden labels (None for new/unlabeled footage). Build one from the DB
171     with ``load_video_candidates``, or directly in tests.
172     """
173
174     name: str
175     intervals: List[Interval]
176     features: Dict[str, List[float]] = field(default_factory=dict)
177     gts: Optional[List[Interval]] = None
178
179
180     def _feature_column(vc: VideoCandidates, name: str) -> List[float]:
181         """Persisted feature column, defaulting missing/short columns to 0.0 (matches
182         ``training.extract_yolo_features``, which lets a sparse/old row still vectorise)."""
183         col = vc.features.get(name)
184         n = len(vc.intervals)
185         if col is None:
186             logger.warning("variant references feature %r absent from %s ? treating as 0.0",
187                            name, vc.name)
188             return [0.0] * n
189         if len(col) != n:
190             raise ExperimentError(
191                 f"feature {name!r} has {len(col)} values but {vc.name} has {n} candidates")
192         return [float(x) for x in col]
193
194
195     def _feature_matrix(vc: VideoCandidates, feature_names: Sequence[str]) ->
196     List[List[float]]:
197         cols = [_feature_column(vc, name) for name in feature_names]
198         return [list(row) for row in zip(*cols)] if cols else [[] for _ in vc.intervals]

```

```

199
200 def _gate_mask(variant: Variant, vc: VideoCandidates) -> List[bool]:
201     """Boolean keep-mask from the decision gates (Gate A net-crossings, Gate B
202     players)."""
203     n = len(vc.intervals)
204     mask = [True] * n
205     if variant.min_crossings > 0:
206         col = _feature_column(vc, "net_crossings")
207         mask = [m and (col[i] >= variant.min_crossings) for i, m in enumerate(mask)]
208     if variant.min_players > 0:
209         col = _feature_column(vc, "players_in_court")
210         mask = [m and (col[i] >= variant.min_players) for i, m in enumerate(mask)]
211     return mask
212
213 def predict(variant: Variant, vc: VideoCandidates,
214             model: Optional[LogisticRegression] = None,
215             threshold: Optional[float] = None,
216             calibrator: Optional[Callable[[float], float]] = None) -> List[Interval]:
217     """Variant prediction: gate by persisted features, optionally filter by a learned
218     model, then merge. Pure and deterministic.
219
220     ``model`` (an already-fitted `LogisticRegression`) is supplied by `compare` per LOVO
221     fold; if omitted and the variant pins ``classifier_model``, that frozen model is
222     loaded. A learned variant with neither a supplied nor a pinned model raises ? there
223     is nothing to score with. ``threshold``/``calibrator`` (both fitted on the TRAINING
224     fold) override the variant defaults when the C2 / threshold-in-LOVO path supplies
225     them.
226     """
227     mask = _gate_mask(variant, vc)
228     if variant.feature_names:
229         if model is None:
230             if variant.classifier_model is None:
231                 raise ExperimentError(
232                     f"variant {variant.name!r} is learned but has no model to predict "
233                     f"with (pass a fitted model or set classifier_model)")
234             model = LogisticRegression.load(variant.classifier_model)
235             proba = [float(p) for p in model.predict_proba(_feature_matrix(vc,
236 variant.feature_names))]
237         if calibrator is not None:
238             proba = [calibrator(p) for p in proba]
239             thr = variant.threshold if threshold is None else threshold
240             mask = [m and (proba[i] >= thr) for i, m in enumerate(mask)]
241         kept = [iv for iv, keep in zip(vc.intervals, mask) if keep]
242         return merge_overlaps(kept, gap_tol=variant.merge_gap)
243
244 def _calibrate_and_tune(variant: Variant, model: LogisticRegression,
245                        train_videos: Sequence[VideoCandidates], iou: float
246                        ) -> "tuple[Optional[Callable[[float], float]],
247 Optional[float]]":
248     """Fit a calibrator and/or pick the operating threshold on the TRAINING fold only
249     (C2 + threshold-in-LOVO). Returns ``(calibrator, threshold)`` ? either may be None
250     (use the variant default). Honest: never looks at the held-out video.
251     """
252     calibrator: Optional[Callable[[float], float]] = None
253     if variant.calibrate:
254         raw: List[float] = []
255         y: List[int] = []
256         for vc in train_videos:
257             raw.extend(float(p) for p in
258 model.predict_proba(_feature_matrix(vc, variant.feature_names or
259 [])))
260             y.extend(label_candidates(vc.intervals, vc.gts or [], iou))
261         if variant.calibrate == "platt":

```

```

259         calibrator = fit_platt(raw, y)
260     elif variant.calibrate == "isotonic":
261         calibrator = fit_isotonic(raw, y)
262     else:
263         raise ExperimentError(f"unknown calibrate={variant.calibrate!r} (use
'platt'/'isotonic')")
264     threshold: Optional[float] = None
265     if variant.tune_threshold:
266         best_t, best_f1 = variant.threshold, -1.0
267         for k in range(1, 20):             # grid 0.05..0.95 on the train
fold's rally F1
268             t = round(0.05 * k, 2)
269             f1s = [score(predict(variant, vc, model=model, threshold=t,
calibrator=calibrator),
270                        vc.gts or [], iou).f1 for vc in train_videos]
271             mean_f1 = sum(f1s) / len(f1s) if f1s else 0.0
272             if mean_f1 > best_f1:
273                 best_f1, best_t = mean_f1, t
274             threshold = best_t
275     return calibrator, threshold
276
277
278     # ----- variant scoring
279
280
281     def _train(variant: Variant, videos: Sequence[VideoCandidates], iou: float) ->
LogisticRegression:
282         """Fit the variant's classifier on the pooled candidates of ``videos`` (labels via
the
283         evaluator's own IoU matcher). The honest-LOVO training step from `distill_local`. """
284         X: List[List[float]] = []
285         y: List[int] = []
286         for vc in videos:
287             if vc.gts is None:
288                 raise ExperimentError(f"cannot train: {vc.name} has no golden labels")
289             X.extend(_feature_matrix(vc, variant.feature_names or []))
290             y.extend(label_candidates(vc.intervals, vc.gts, iou))
291         if not X:
292             raise ExperimentError("cannot train a learned variant on zero candidates")
293         return LogisticRegression().fit(X, y, variant.feature_names)
294
295
296     def evaluate_variant(videos: Sequence[VideoCandidates], variant: Variant,
297                        iou: float = 0.5, honest: bool = True) -> Dict[str, Any]:
298         """Score a variant on a labeled corpus. Returns per-video Scores + predictions + a
299         mean aggregate.
300
301         Prediction policy:
302         - **static variant** (no learned filter): predict each video directly ? no
training,
303         so no leakage; per-video scoring is already honest.
304         - **learned, pinned model** (`classifier_model` set): score every video with the
305         SHIPPED model. ? optimistic if those videos were in its training set ? use this
to
306         report "how the shipped artifact does here", not for the promotion decision.
307         - **learned recipe** (`feature_names`, no pinned model) + `honest=True`: LOVO
?
308         train on the OTHER videos, predict the held-out one. The number to trust.
309         - learned recipe + `honest=False`: train on all, predict each (overfit; sanity
only).
310         """
311         for vc in videos:
312             if vc.gts is None:
313                 raise ExperimentError(f"{vc.name} has no golden labels; use
compare_unlabeled")

```

```

314
315     per_video: List[Dict[str, Any]] = []
316     predictions: Dict[str, List[Interval]] = {}
317
318     recipe_lovo = bool(variant.feature_names) and variant.classifier_model is None
319     pinned_model = (LogisticRegression.load(variant.classifier_model)
320                     if variant.classifier_model is not None else None)
321     all_model = None
322     if recipe_lovo and not honest:
323         all_model = _train(variant, videos, iou)
324
325     for i, vc in enumerate(videos):
326         cal: Optional[Callable[[float], float]] = None
327         thr: Optional[float] = None
328         if recipe_lovo and honest:
329             others = [v for j, v in enumerate(videos) if j != i]
330             if not others:
331                 raise ExperimentError("LOVO needs >=2 videos; pass honest=False or a
pinned model")
332             model: Optional[LogisticRegression] = _train(variant, others, iou)
333             if variant.tune_threshold or variant.calibrate: # operating point from
the train fold
334                 assert model is not None # _train never returns
None
335                 cal, thr = _calibrate_and_tune(variant, model, others, iou)
336             elif recipe_lovo: # honest=False ? one model over all
videos
337                 model = all_model
338             else: # static variant or pinned model
339                 model = pinned_model
340             preds = predict(variant, vc, model=model, threshold=thr, calibrator=cal)
341             assert vc.gts is not None # guarded at function top
342             sc = score(preds, vc.gts, iou)
343             per_video.append({"video": vc.name, "num_pred": len(preds), **{m: getattr(sc, m)
for m in _METRICS}})
344             predictions[vc.name] = preds
345
346         aggregate = {m: (sum(p[m] for p in per_video) / len(per_video) if per_video else
0.0)
347                     for m in _METRICS}
348     return {
349         "variant": variant.name,
350         "spec_hash": variant.spec_hash(),
351         "is_learned": variant.is_learned(),
352         "honest_lovo": recipe_lovo and honest,
353         "per_video": per_video,
354         "aggregate": aggregate,
355         "predictions": predictions,
356     }
357
358
359     def _ci_block(eval_v: Dict[str, Any]) -> Dict[str, Any]:
360         """Per-metric mean + bootstrap CI over the per-video scores (C1 ? never a bare
mean)."""
361         return {m: _stats.mean_with_ci([p[m] for p in eval_v["per_video"]]) for m in
_METRICS}
362
363
364     def _segmentation_block(videos: Sequence[VideoCandidates], eval_v: Dict[str, Any]) ->
Dict[str, Any]:
365         """Mean F1@k + segment-count ratio across videos (C4 ? the over-segmentation tIoU
hides)."""
366         gts_by_name = {v.name: (v.gts or []) for v in videos}
367         overlaps = _segm.DEFAULT_OVERLAPS
368         f1k: Dict[float, List[float]] = {ov: [] for ov in overlaps}

```

```

369 ratios: List[float] = []
370 for name, preds in eval_v.get("predictions", {}).items():
371     gts = gts_by_name.get(name, [])
372     got = _segm.f1_at_overlaps(preds, gts, overlaps)
373     for ov in overlaps:
374         f1k[ov].append(got[ov])
375     ratios.append(_segm.segment_count_ratio(preds, gts))
376
377 def _mean(xs: List[float]) -> float:
378     return sum(xs) / len(xs) if xs else 0.0
379 out: Dict[str, Any] = {f"f1@{int(ov * 100)}": _mean(vals) for ov, vals in
f1k.items()}
380 out["segment_count_ratio"] = _mean(ratios)
381 return out
382
383
384 def honest_summary(videos: Sequence[VideoCandidates], eval_a: Dict[str, Any],
385                   eval_b: Dict[str, Any]) -> Dict[str, Any]:
386     """C1 + C4 honest reporting layered over a compare() pair (additive,
EXPERIMENT_HARNESS $6).
387
388     ``per_video`` lists are video-aligned (``evaluate_variant`` iterates ``videos`` in
order),
389     so ``paired_delta_f1`` pairs B?A by video ? the promotion-grade view: a lift is real
when
390     the ?F1 CI clears 0 *and* the sign test agrees, not when a bare mean ticked up.
391     """
392     a_f1 = [p["f1"] for p in eval_a["per_video"]]
393     b_f1 = [p["f1"] for p in eval_b["per_video"]]
394     return {
395         "variant_a_ci": _ci_block(eval_a),
396         "variant_b_ci": _ci_block(eval_b),
397         "paired_delta_f1": _stats.paired_delta_ci(a_f1, b_f1),
398         "segmentation": {"variant_a": _segmentation_block(videos, eval_a),
399                         "variant_b": _segmentation_block(videos, eval_b)},
400     }
401
402
403 def compare(videos: Sequence[VideoCandidates], variant_a: Variant, variant_b: Variant,
404            iou: float = 0.5, honest: bool = True, honest_stats: bool = True) ->
Dict[str, Any]:
405     """Offline labeled A/B (``EXPERIMENT_HARNESS.md`` $5.1). Scores both variants on the
same
406     golden corpus and reports the **B ? A** deltas, per video and in aggregate.
407
408     The deltas use the existing ``metrics.score`` (per video) +
``calibration.pct_improvement``;
409     learned variants are scored LOVO-honestly. The result is a self-contained record fit
for
410     ``record_experiment`` ? variant specs + hashes + the label-set fingerprint travel with
it.
411
412     ``honest_stats`` (default True) **adds** a ``"honest"`` block ? per-metric bootstrap
CIs, a
413     paired ?F1 CI + exact sign test (C1), and F1@k + segment-count ratio (C4) ?
*alongside* the
414     existing bare-mean fields (additive: nothing existing changes). Set False for the
legacy
415     shape. This is the measurement-honesty wiring (``ANALYZERS_AND_RECONCILERS.md``
$13/C1+C4): a
416     bare LOVO mean at n?6 is not trustworthy on its own.
417     """
418     if len(videos) < 1:
419         raise ExperimentError("compare needs at least one labeled video")
420     eval_a = evaluate_variant(videos, variant_a, iou, honest)

```

```

421     eval_b = evaluate_variant(videos, variant_b, iou, honest)
422
423     delta = {m: eval_b["aggregate"][m] - eval_a["aggregate"][m] for m in _METRICS}
424     delta_pct = {m: pct_improvement(eval_a["aggregate"][m], eval_b["aggregate"][m]) for
m in _METRICS}
425
426     by_video_a = {p["video"]: p for p in eval_a["per_video"]}
427     per_video_delta = [
428         {"video": p["video"], **{m: p[m] - by_video_a[p["video"]][m] for m in _METRICS}}
429         for p in eval_b["per_video"]
430     ]
431
432     # One fingerprint over the whole label set ? detects "the deltas were measured
against
433     # different labels" the same way regression.gt_hash guards the no-regression
baseline.
434     all_gts: List[Interval] = [iv for vc in videos for iv in (vc.gts or [])]
435
436     result: Dict[str, Any] = {
437         "iou": iou,
438         "label_set_hash": gt_hash(all_gts),
439         "num_videos": len(videos),
440         "variant_a": eval_a,
441         "variant_b": eval_b,
442         "delta": delta,
443         "delta_pct": delta_pct,
444         "per_video_delta": per_video_delta,
445     }
446     if honest_stats:
447         result["honest"] = honest_summary(videos, eval_a, eval_b)
448     return result
449
450
451     # ----- SxS on unlabeled video
452
453
454     def compare_unlabeled(video: VideoCandidates, variant_a: Variant, variant_b: Variant,
455         agree_iou: float = 0.5) -> Dict[str, Any]:
456         """Side-by-side on NEW footage with no ground truth (`EXPERIMENT_HARNESS.md` §5.2).
457
458         With no labels there is no lift to measure ? instead this surfaces where the two
459         variants **agree vs diverge**, which is exactly what a human reviews (and what two
460         highlight reels would show side by side):
461
462         - ``agreed`` ? windows both variants found (matched at ``agree_iou``);
463         - ``a_only`` ? A fired, B suppressed (B's added precision, if A was over-firing);
464         - ``b_only`` ? B fired, A did not (B's new detections ? real recall or new FPs:
465             the human / a later golden label adjudicates).
466
467         Both variants must be predictable without labels: a learned variant therefore needs
a
468         pinned ``classifier_model`` (you cannot train on an unlabeled video). Reuses
469         ``metrics.match_intervals`` ? no new matching logic.
470         """
471         preds_a = predict(variant_a, video)
472         preds_b = predict(variant_b, video)
473         mr = match_intervals(preds_a, preds_b, agree_iou)
474
475         agreed = [{"a": preds_a[m.pred_index], "b": preds_b[m.gt_index], "iou": m.iou}
476             for m in mr.matches]
477         agree_iou = [m.iou for m in mr.matches]
478         a_only = [preds_a[i] for i in mr.unmatched_pred_indices]
479         b_only = [preds_b[i] for i in mr.unmatched_gt_indices]
480
481     def _dur(ivs: List[Interval]) -> float:

```

```

482         return sum(e - s for s, e in ivs)
483
484     return {
485         "video": video.name,
486         "agree_iou": agree_iou,
487         "variant_a": variant_a.name,
488         "variant_b": variant_b.name,
489         "num_a": len(preds_a),
490         "num_b": len(preds_b),
491         "num_agreed": len(agreed),
492         "mean_agree_iou": (sum(agree_ious) / len(agree_ious)) if agree_ious else 0.0,
493         "duration_a": _dur(preds_a),
494         "duration_b": _dur(preds_b),
495         "agreed": agreed,
496         "a_only": a_only,
497         "b_only": b_only,
498     }
499
500
501 # ----- leaderboard / DB
502
503
504 def record_experiment(result: Dict[str, Any], path: str = DEFAULT_LEADERBOARD_PATH,
505                     timestamp: Optional[str] = None) -> Dict[str, Any]:
506     """Append a compact, reproducible record of a `compare()` result to the JSON
507     leaderboard (`EXPERIMENT_HARNESS.md` §6: experiments are records). Generalises
508     `regression_baseline.json`; deliberately the size of a baseline file, not a service
509     (§9). Returns the row written. ``timestamp`` is injectable for deterministic tests.
510     """
511     row: Dict[str, Any] = {
512         "timestamp": timestamp or datetime.now(timezone.utc).isoformat(),
513         "iou": result.get("iou"),
514         "label_set_hash": result.get("label_set_hash"),
515         "num_videos": result.get("num_videos"),
516         "variant_a": {"name": result["variant_a"]["variant"],
517                      "spec_hash": result["variant_a"]["spec_hash"],
518                      "aggregate": result["variant_a"]["aggregate"]},
519         "variant_b": {"name": result["variant_b"]["variant"],
520                      "spec_hash": result["variant_b"]["spec_hash"],
521                      "aggregate": result["variant_b"]["aggregate"]},
522         "delta": result.get("delta"),
523     }
524     # Record the honest ?F1 (CI + sign test) when present, so the leaderboard carries
525     the
526     # promotion-grade evidence, not just the bare-mean delta (additive ? C1).
527     honest = result.get("honest") or {}
528     if honest.get("paired_delta_f1"):
529         row["honest_delta_f1"] = honest["paired_delta_f1"]
530     entries: List[Dict[str, Any]] = []
531     if os.path.isfile(path):
532         with open(path, encoding="utf-8") as f:
533             entries = json.load(f)
534     entries.append(row)
535     os.makedirs(os.path.dirname(path) or ".", exist_ok=True)
536     with open(path, "w", encoding="utf-8") as f:
537         json.dump(entries, f, indent=2)
538     logger.info("recorded experiment %s vs %s (?f1=%0.3f) -> %s",
539               row["variant_a"]["name"], row["variant_b"]["name"],
540               (result.get("delta") or {}).get("f1", 0.0), path)
541     return row
542
543 def load_video_candidates(db: Any, video_id: str, *, detector: Optional[str] = None,
544                         gts: Optional[List[Interval]] = None,
545                         feature_names: Sequence[str] = ()) -> VideoCandidates:

```



```

546         """Bridge persisted detection ? `VideoCandidates` via
`Database.query_candidate_segments`.
547
548         ``feature_names`` selects which persisted features to lift out of each row's
``stats``
549         JSON (``"duration"`` is derived from the window bounds). This is the only DB-
touching
550         function; the comparison itself is pure and offline. Exercised owner-side on a real
DB,
551         but works against any `Database` (including an in-memory one in tests).
552         """
553         rows = db.query_candidate_segments(video_id, detector=detector)
554         intervals: List[Interval] = [(float(r["start_time"]), float(r["end_time"])) for r in
rows]
555         features: Dict[str, List[float]] = {}
556         for name in feature_names:
557             if name == "duration":
558                 features[name] = [float(r["end_time"]) - float(r["start_time"]) for r in
rows]
559             else:
560                 features[name] = [float((r.get("stats") or {}).get(name, 0.0)) for r in
rows]
561         return VideoCandidates(name=video_id, intervals=intervals, features=features,
gts=gts)
562

```

11. assistant (2026-06-14T09:49:49.624Z)

Now I understand the harness architecture. Let me inspect a sample golden JSON to see which feature columns exist (shuttle/person, and whether audio is present yet), the feature name lists, the eval_stats module, and the existing baselines/manifest.

12. user (2026-06-14T09:49:54.081Z)

```

adarsh_avi_singles          n_cand=172 n_pos=80 n_gts=96 keys=['start', 'end', 'shuttle',
'person', 'label', 'audio'] audio=True
    shuttle feats: ['duration', 'peak_velocity', 'mean_velocity', 'active_ratio',
'net_crossings']
    person  feats: ['players_in_court', 'two_sided_motion', 'mean_player_motion',
'in_court_ratio', 'shuttle_player_colocation']
    audio   feats: ['audio_hit_count', 'audio_hit_rate', 'audio_hit_strength_mean']

```

13. user (2026-06-14T09:49:54.356Z)

```

1         """Honest small-N evaluation statistics (course-corrections C1 + C3).
2
3         The ~6-match golden corpus + a regularised logistic arbiter is exactly the regime where
a
4         single leave-one-video-out *mean* lies: leave-one-out anti-correlates each fold's
training
5         and held-out label means at small N, and one lucky/unlucky fold swings the headline. The
6         ``compare``/LOVO machinery already produces **per-video held-out scores** ? this module
7         turns those into an honest summary:
8
9         - **C1 ? never a single mean.** ``mean_with_ci`` / ``bootstrap_ci`` report a confidence
10        interval; ``paired_sign_test`` and ``paired_delta_ci`` grade a B?A improvement by how
11        many *videos* it wins on (the distribution-free test the repo already leans on:
12        "sign test needs n ? ~n=10/9-wins for p<0.05"), so the promotion gate can't promote
noise.
13        - **C3 ? honest stacking.** ``out_of_fold_predictions`` / ``grouped_folds`` give a
learned
14        producer's *out-of-fold* outputs so feeding them to the arbiter doesn't leak (group by
15        match, never a held-out window from the same video).
16

```

```

17 See ``docs/ANALYZERS_AND_RECONCILERS.md`` §13/C1+C3. All **additive**, pure (stdlib
only),
18 deterministic (bootstrap takes an explicit ``seed``) ? nothing here changes existing
19 ``compare``/``calibration.leave_one_video_out`` behaviour; callers opt in.
20 """
21 from __future__ import annotations
22
23 import math
24 import random
25 import statistics
26 from typing import Any, Callable, Dict, List, Optional, Sequence, Tuple
27
28 Interval = Tuple[float, float]
29 FitPredict = Callable[[List[int], List[int]], List[Any]]
30
31
32 def bootstrap_ci(values: Sequence[float], confidence: float = 0.95,
33                 n_boot: int = 2000, seed: int = 0) -> Tuple[float, float]:
34     """Percentile bootstrap CI for the mean of ``values``. Deterministic given ``seed``.
35
36     ``n<=1`` returns a degenerate ``(v, v)`` / ``(0, 0)`` interval ? there is no spread
to
37     resample. With n=6 (our corpus) this is the honest "how wide could the mean be?"
band.
38     """
39     vals = [float(v) for v in values]
40     n = len(vals)
41     if n == 0:
42         return (0.0, 0.0)
43     if n == 1:
44         return (vals[0], vals[0])
45     rng = random.Random(seed)
46     means: List[float] = []
47     for _ in range(max(1, n_boot)):
48         means.append(sum(vals[rng.randrange(n)] for _ in range(n)) / n)
49     means.sort()
50     alpha = (1.0 - confidence) / 2.0
51     last = len(means) - 1
52     lo = means[max(0, int(math.floor(alpha * last)))]
53     hi = means[min(last, int(math.ceil((1.0 - alpha) * last)))]
54     return (lo, hi)
55
56
57 def mean_with_ci(values: Sequence[float], confidence: float = 0.95,
58                 n_boot: int = 2000, seed: int = 0) -> Dict[str, float]:
59     """``{mean, std, n, ci_low, ci_high, confidence}`` ? report this, never a bare mean
(C1)."""
60     vals = [float(v) for v in values]
61     n = len(vals)
62     mean = statistics.mean(vals) if vals else 0.0
63     std = statistics.pstdev(vals) if n > 1 else 0.0
64     lo, hi = bootstrap_ci(vals, confidence, n_boot, seed)
65     return {"mean": mean, "std": std, "n": float(n),
66           "ci_low": lo, "ci_high": hi, "confidence": confidence}
67
68
69 def paired_sign_test(deltas: Sequence[float], eps: float = 0.0) -> Dict[str, float]:
70     """Two-sided exact sign test over per-fold deltas (B ? A).
71
72     ``wins`` = folds where B beat A by more than ``eps``; ``losses`` the reverse; ties
73     excluded. ``p_value`` is the exact two-sided binomial tail under p=0.5 ? the
74     distribution-free "did B win on enough *videos*?" check.
75     """
76     wins = sum(1 for d in deltas if d > eps)
77     losses = sum(1 for d in deltas if d < -eps)

```

```

78     ties = len(deltas) - wins - losses
79     n_eff = wins + losses
80     if n_eff == 0:
81         p = 1.0
82     else:
83         k = min(wins, losses)
84         tail = sum(math.comb(n_eff, j) for j in range(k + 1)) * (0.5 ** n_eff)
85         p = min(1.0, 2.0 * tail)
86     return {"wins": float(wins), "losses": float(losses), "ties": float(ties),
87            "n_effective": float(n_eff), "p_value": p}
88
89
90 def paired_delta_ci(a_values: Sequence[float], b_values: Sequence[float],
91                    confidence: float = 0.95, n_boot: int = 2000, seed: int = 0,
92                    eps: float = 0.0) -> Dict[str, Any]:
93     """Paired (by fold) B ? A summary: mean delta + bootstrap CI + the sign test (C1).
94
95     ``a_values``/``b_values`` are per-fold held-out scores for variants A and B, aligned
96     by
97     fold (same video order). A promotion is honest when the delta CI clears 0 *and* the
98     sign
99     test agrees ? not when a single mean ticked up.
100     """
101     n = min(len(a_values), len(b_values))
102     deltas = [float(b_values[i]) - float(a_values[i]) for i in range(n)]
103     summary = mean_with_ci(deltas, confidence, n_boot, seed)
104     return {
105         "mean_delta": summary["mean"],
106         "ci_low": summary["ci_low"],
107         "ci_high": summary["ci_high"],
108         "n": float(n),
109         "ci_excludes_zero": (summary["ci_low"] > 0.0) or (summary["ci_high"] < 0.0),
110         "sign_test": paired_sign_test(deltas, eps),
111     }
112
113 def grouped_folds(groups: Sequence[Any]) -> List[Tuple[List[int], List[int]]]:
114     """Leave-one-group-out splits ``(train_idx, test_idx)`` ? group by match, the honest
115     unit.
116
117     The field's name for "LOVO by match" is leave-one-group-out (LOGO). Group order is
118     first-seen (deterministic). Groups with no complement (a single group) yield no
119     fold.
120
121     """
122     seen: List[Any] = []
123     for g in groups:
124         if g not in seen:
125             seen.append(g)
126     n = len(groups)
127     folds: List[Tuple[List[int], List[int]]] = []
128     for g in seen:
129         test = [i for i in range(n) if groups[i] == g]
130         train = [i for i in range(n) if groups[i] != g]
131         if test and train:
132             folds.append((train, test))
133     return folds
134
135 def out_of_fold_predictions(groups: Sequence[Any], fit_predict: FitPredict) ->
136 List[Optional[Any]]:
137     """Out-of-fold predictions for honest stacking (C3).
138
139     For each leave-one-group-out fold, ``fit_predict(train_idx, test_idx)`` trains on
140     the
141     other groups and returns predictions for ``test_idx`` (in that order); results are

```

```

137         scattered back to a list aligned to the input items. So a *learned* producer's
outputs,
138         when used as arbiter features, never include in-sample predictions ? the leakage
path
139         that inflates measured lift on a tiny corpus. Items in a singleton group stay
``None``
140         (needs >= 2 groups to predict).
141         """
142         n = len(groups)
143         oof: List[Optional[Any]] = [None] * n
144         for train_idx, test_idx in grouped_folds(groups):
145             preds = fit_predict(train_idx, test_idx)
146             for pos, ti in enumerate(test_idx):
147                 if pos < len(preds):
148                     oof[ti] = preds[pos]
149         return oof
150

```

14. assistant (2026-06-14T09:50:00.738Z)

Audio columns already exist in the JSONs ? no ``--augment`` needed. Now let me confirm all 6 JSONs have audio, check the manifest, existing baselines, and QUALITY_ITERATIONS for the known numbers.

15. user (2026-06-14T09:50:03.902Z)

```

adarsh_avi_singles          n_cand=172 n_pos= 80 n_gts=96 audio=True fps=59.94
gbaaddy                    n_cand= 82 n_pos= 15 n_gts=48 audio=True fps=59.94
kushagra_singles           n_cand= 20 n_pos=  4 n_gts=32 audio=True fps=59.94
largetest_doubles          n_cand= 86 n_pos= 40 n_gts=49 audio=True fps=29.97
mahadevpura_singles        n_cand= 57 n_pos= 20 n_gts=31 audio=True fps=29.97
testlarge_short            n_cand= 14 n_pos=  3 n_gts= 6 audio=True fps=29.97
---MANIFEST---
[
  {
    "name": "adarsh_avi_singles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/adarsh_avi_full_traj.csv",
    "fps": 59.94,
    "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden Labelled/Adarsh and
Avi.rallies.csv"
  },
  {
    "name": "kushagra_singles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/kushagra_traj.csv",
    "fps": 59.94,
    "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/KushagraYashAviFirstVideo.rallies.csv"
  },
  {
    "name": "largetest_doubles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/largetestvideo_traj.csv",
    "fps": 29.97,
    "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/LargeTestVideo.rallies.csv"
  },
  {
    "name": "gbaaddy",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/gbaaddy_traj.csv",

```

```

    "fps": 59.94,
    "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/gBaaddy20May2026.rallies.csv"
  },
  {
    "name": "testlarge_short",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/testlarge_traj.csv",
    "fps": 29.97,
    "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/TestLargeVideo.rallies.csv"
  },
  {
    "name": "mahadevpura_singles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/mahadevpura_traj.csv",
    "fps": 29.97,
    "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/MahadevpuraSingles.rallies.csv"
  }
]

```

16. user (2026-06-14T09:50:05.327Z)

```

experiment_leaderboard.json
gemini_wrapper_2026-06-10.json
heuristic_lovo_n6.json
---LEADERBOARD---
rows: 2
{
  "timestamp": "2026-06-14T06:23:51.507393+00:00",
  "iou": 0.5,
  "label_set_hash": "50d98ffbc370",
  "num_videos": 6,
  "variant_a": {
    "name": "shuttle_only",
    "spec_hash": "47c867c36f2f",
    "aggregate": {
      "precision": 0.31807603780927995,
      "recall": 0.3101947096042718,
      "f1": 0.3068611057861412,
      "mean_iou": 0.5791831889148886
    }
  },
  "variant_b": {
    "name": "shuttle_person_fusion",
    "spec_hash": "253f4709b4c4",
    "aggregate": {
      "precision": 0.3617228001510626,
      "recall": 0.2708939086387243,
      "f1": 0.2863398846047612,
      "mean_iou": 0.5690111885714347
    }
  },
  "delta": {
    "precision": 0.04364676234178266,
    "recall": -0.0393008009655475,
    "f1": -0.02052122118138,
    "mean_iou": -0.01017200034345389
  },
  "honest_delta_f1": {
    "mean_delta": -0.02052122118137999,

```

```

"ci_low": -0.16462806145542527,
"ci_high": 0.16224446866060574,
"n": 6.0,
"ci_excludes_zero": false,
"sign_test": {
  "wins": 2.0,
  "losses": 4.0,
  "ties": 0.0,
  "n_effective": 6.0,
  "p_value": 0.6875
}
}
}
{
  "timestamp": "2026-06-14T06:24:56.487711+00:00",
  "iou": 0.5,
  "label_set_hash": "50d98ffbc370",
  "num_videos": 6,
  "variant_a": {
    "name": "shuttle_person",
    "spec_hash": "ee89f1e41604",
    "aggregate": {
      "precision": 0.3617228001510626,
      "recall": 0.2708939086387243,
      "f1": 0.2863398846047612,
      "mean_iou": 0.5690111885714347
    }
  },
  "variant_b": {
    "name": "shuttle_person_audio",
    "spec_hash": "d68c39eba87c",
    "aggregate": {
      "precision": 0.4530798177857001,
      "recall": 0.3408138120839734,
      "f1": 0.3655826552797436,
      "mean_iou": 0.7074802852140344
    }
  },
  "delta": {
    "precision": 0.0913570176346375,
    "recall": 0.06991990344524907,
    "f1": 0.07924277067498242,
    "mean_iou": 0.13846909664259965
  },
  "honest_delta_f1": {
    "mean_delta": 0.07924277067498244,
    "ci_low": -0.053508315003560535,
    "ci_high": 0.2029045239922802,
    "n": 6.0,
    "ci_excludes_zero": false,
    "sign_test": {
      "wins": 4.0,
      "losses": 2.0,
      "ties": 0.0,
      "n_effective": 6.0,
      "p_value": 0.6875
    }
  }
}
}

```

17. assistant (2026-06-14T09:50:09.187Z)

All 6 JSONs have audio columns ? no `--augment` needed. The leaderboard already has P3 (person) and P4 (incremental audio). Now let me read the existing baselines and the QUALITY_ITERATIONS doc for the known Gate-A numbers and the heuristic/Gen-0/wrapper baselines.

18. user (2026-06-14T09:50:11.977Z)

=== heuristic_lovo_n6.json ===

```
{
  "name": "heuristic_velocity_windowing_LOV0",
  "description": "The honest cross-video floor the Gen-0 learned model must beat: calibrated velocity-windowing heuristic, leave-one-video-out over the n=6 owned golden corpus.",
  "computed": "2026-06-10, reproduced via: python -m backend.eval.calibrate_local --manifest output/human_lovo_manifest.json --metric f1 (matches docs/archives/quality-iterations/2026-06-multivideo-lovo)",
  "mean_f1": 0.48,
  "std": 0.242,
  "per_video": {
    "adarsh_avi_singles": 0.751,
    "kushagra_singles": 0.157,
    "largetest_doubles": 0.737,
    "gbaaddy": 0.309,
    "testlarge_short": 0.267,
    "mahadevpura_singles": 0.657
  }
}
```

=== gemini_wrapper_2026-06-10.json ===

```
{
  "name": "gemini_wrapper_baseline",
  "description": "Commodity-VLM wrapper baseline: the full pipeline's `gemini` segmenter scored vs human golden labels at IoU 0.5. The commoditization/moat reference the owned model must be judged against \u2014 on multi-court footage the wrapper drops 0.93 -> 0.66, so the owned model's target is the multi-court number, not the clean-footage one.",
  "model": "gemini-flash-latest (resolved 2026-06-10; 3.5-flash era)",
  "conditions": "serial uploads (ai_max_workers=1), chunks re-encoded to 1280px w/ audio (ai_chunk_scale_width), 6-attempt exponential generate retry; min_segment_duration=3.0 filter active (rejects real ~2.5s rallies \u2014 see notes)",
  "per_video": {
    "mahadevpura_singles": {
      "f1": 0.93,
      "precision": 1.0,
      "recall": 0.87,
      "mean_iou": 0.8,
      "boundary_err_s": 0.7,
      "n_gt": 31,
      "n_pred": 27,
      "notes": "3 of 4 FNs were detected by Gemini but rejected by min_segment_duration=3.0 (2.5s rallies). Chunk 6 (unlabeled tail) lost to 503s \u2014 no effect on the labeled window."
    },
    "kushagra_singles": {
      "f1": 0.66,
      "precision": 0.69,
      "recall": 0.62,
      "mean_iou": 0.73,
      "boundary_err_s": 1.2,
      "n_gt": 32,
      "n_pred": 29,
      "notes": "Multi-court (contrast 1.4x). Full 7/7 chunk coverage, audio intact. FPs cluster as background-game pickups + one 53s dead-time merge \u2014 the exact confound the contrast metric predicts."
    }
  },
  "comparison_at_time_of_run": {
    "heuristic_lovo": {
      "mahadevpura_singles": 0.657,
      "kushagra_singles": 0.157
    },
    "gen0_learned_lovo": {
      "mahadevpura_singles": 0.744,

```

```

    "kushagra_singles": 0.561
  }
},
"reliability_note": "gemini-flash-latest served 503 capacity waves throughout (one chunk needed
a ~5min backoff); a sole-provider dependency is operationally flaky \u2014 the provider-registry
fallback design is load-bearing.",
"cost_note": "well under $1 for both videos at flash pricing (serial, ~14 chunk requests
total).".
}

```

19. user (2026-06-14T09:50:12.798Z)

```

eval_baselines/experiment_leaderboard.json
eval_baselines/gemini_wrapper_2026-06-10.json
eval_baselines/heuristic_lovo_n6.json
---
docs/QUALITY_ITERATIONS.md
212 docs/QUALITY_ITERATIONS.md

```

20. user (2026-06-14T09:50:15.304Z)

```

1      # Rally-detection quality iterations ? living log
2
3      **Purpose.** A running, reverse-chronological log of how we evolve rally-detection
quality: each
4      iteration's hypothesis, what shipped, the measured (or expected) effect, and the lesson.
This is the
5      **live** continuation of the historical trail in
6      [`archives/quality-iterations/`](archives/quality-iterations/README.md) ? per the
archive policy
7      (`CLAUDE.md`), **terminal-state** (DONE/REJECTED) work gets a dated subdir there; this
file stays at
8      `docs/` top level and keeps growing, then its entries graduate to the archive when they
conclude.
9
10     **Methodology (non-negotiable).** Every quality change rides the experiment harness
11     ([EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)): new cues land **flag-gated, default
OFF**;; they're
12     graded **offline against golden labels** (per-video + LOVO-honest); promotion to default
happens **only
13     on measured lift** through the `run_regression.py` gate; **rollback = flip a config
flag, never
14     `git revert`**. Cues are designed in
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md); scoring
15     in [EVALUATION.md](EVALUATION.md); the 2-layer model in
16     [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md).
17
18     **Golden corpus = the 6 labelled matches** in `output/human_lovo_manifest.json` (owner-
confirmed
19     2026-06-13: use all 6 ? `adarsh_avi_singles, kushagra_singles, largetest_doubles,
gbaaddy,
20     testlarge_short, mahadevpura_singles`). The model trained on these is what we run
alongside the existing
21     detector on NEW footage for the side-by-side review.
22
23     ---
24
25     ## Baselines that drive strategy (IoU 0.5 vs human golden labels)
26
27     | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |
28     |---|---|---|
29     | Heuristic (velocity windowing) | 0.657 | 0.157 |
30     | Gen-0 owned head (LOVO, n=6) | 0.744 | 0.561 |
31     | Two-stage WASB+Gemini refine | 0.83 | ? |
32     | Gemini wrapper (`gemini-flash-latest`) | 0.93 | 0.66 |

```



```

33
34     **Reading:** clean footage is commoditized (serve via Gemini); **multi-court drops the
commodity wrapper
35     to 0.66 ? that is the owned model's ship target.** Baselines tracked in
`eval_baselines/`. The owner's
36     observed false positive ? a player **holding/flicking** the shuttle between points
counted as a rally ?
37     is the precision bug the fusion gates target.
38
39     ---
40
41     ## Current track (live) ? newest first
42
43     ### Fusion P3 + P4 MEASURED on the n=6 golden corpus ? person = no lift, audio =
promising (not promotable) *(2026-06-14)*
44     The GPU person-feature extraction finished all 6 golden videos (`fusion_golden.py` ?
`output/*_golden_features.json`, the
45     replayable substrate); the CPU LOVO drivers then measured the lift with **honest stats
(C1 ? bootstrap 95% CI + paired sign
46     test; never a bare mean at n=6)**:
47
48     | Variant (honest LOVO, IoU ? 0.5) | F1 | ?F1 vs prior | 95% CI | sign test | CI clears
0 |
49     |---|---|---|---|---|---|
50     | shuttle-only | 0.307 | ? | ? | ? | ? |
51     | + person (P3) | 0.286 | **?0.021** (?6.7%) | [?0.165, +0.162] | 2W/4L (p=0.688) |
**No** |
52     | + person + audio (P4 incr.) | 0.366 | **+0.079** (+27.7%) | [?0.054, +0.203] | 4W/2L
(p=0.688) | **No** |
53
54     **Read honestly:** neither cue clears the promotion bar (both CIs span 0, both sign
tests n.s.) at n=6 ? **both stay
55     default-OFF** (promote-only-on-lift). **Person** adds noise (no court polygons ?
`players_in_court` counts spectators);
56     per-video it's a wash (testlarge +0.40, gbaaddy ?0.22). **Audio** is the more promising
lead ? +0.079 / +27.7%, wins 4/6
57     (largetest +0.27, mahadevpura +0.27), and shuttle+person+audio (0.366) lands **above**
shuttle-only (0.307) ? but n=6 can't
58     confirm it. **Next:** more matches (raise N until the sign test can resolve ~+0.08), and
measure audio on shuttle-only
59     *directly* (without the person handicap). Recorded in
`eval_baselines/experiment_leaderboard.json`. **Negative/uncertain
60     results are kept on purpose ? this is the honest ablation backbone for the paper aim.**
61
62     ### Acting on the first A/B: Gate-A PROMOTED + learned variant FIXED + GX010125
corroboration *(2026-06-14)*
63     **Gate-A (net-crossings ? 2) ? PROMOTED (first cue promoted via the harness).** The
measured win
64     (?F1 +0.074, **6/6 golden matches, sign-test p=0.031**, CI [+0.042,+0.104]; over-seg
1.68?1.01 ? the
65     [first-A/B archive](archives/quality-iterations/2026-06-first-ab-experiment/README.md))
clears the honest
66     bar. **Recommended config:** `indexing.fusion.min_crossings = 2` on the fusion path (the
eval-only Gate A
67     folded into production, MULTI_SIGNAL_FUSION_PLAN §3.1). ? The offline win is on the
recall-oriented *proposal*
68     candidate set; flipping the **live served default** stays owner-gated until a real
fusion-path run confirms it
69     (the production fusion segmenter is P2/unverified on real video) ? so this records the
promotion + rationale,
70     not a silent default flip. **Corroborated on real footage:** offline re-score of
GX010125_1080p's *cached*
71     trajectory (no GPU, isolated `output/GX010125_run/`) ? 45 candidates ? CONTROL 45
rallies vs **Gate-A 27** ;
72     Gate-A dropped **18 low-crossing windows (~1.7 min)** of likely held-shuttle/feed FPs

```

```

(no golden labels for
73     this footage ? divergence, not F1).
74
75     **Learned variant FIXED ? [#130](https://github.com/avidullu/badminton-highlight-
indexer/pull/130).** The
76     first A/B's learned logistic lost to CONTROL and zeroed `testlarge` ? a fixed-0.5-cutoff
artifact. Added
77     additive/default-OFF `Variant.tune_threshold` (operating point chosen on the **train**
fold) + `calibrate`
78     (Platt/isotonic, C2). Validated on the n=6 golden set: **mean F1 0.307 ? 0.423
(+0.116)** , testlarge
79     0.000 ? 0.300, and the learned variant now **beats CONTROL (0.388)** . Threshold-in-LOVO
is the big lever;
80     calibration is available too.
81
82     **Next:** a real fusion-path run to confirm Gate-A served-side; per-video output layout
83     ([PER_VIDEO_OUTPUT_LAYOUT.md](PER_VIDEO_OUTPUT_LAYOUT.md)) so each video's artifacts are
self-contained.
84
85     ### First end-to-end A/B on the golden corpus (n=6) ? **ARCHIVED** *(2026-06-14)*
86     First real use of the harness + C1?C4 helpers, 100% offline (re-scored persisted WASB
trajectories, no GPU).
87     **Gate-A (net-crossings ? 2) is a measured win** ? F1 +0.074, 6/6 matches, sign-test
**p = 0.031** , CI
88     [+0.042, +0.104], over-segmentation (segment-count ratio) **1.68 ? 1.01** . The window-
level 5-feature logistic
89     **did not clear the bar** (?0.081, 2W/4L; a threshold/calibration artifact ? *not* the
stronger per-frame
90     prototype). The framework held up: **C1 promoted one variant and refused another; C4
surfaced over-segmentation
91     tIoU hid.** Full record + tables ?
92     [archives/quality-iterations/2026-06-first-ab-experiment/](archives/quality-
iterations/2026-06-first-ab-experiment/README.md).
93     Next: Gate-A is a candidate for the no-regression promotion gate; the logistic failure
is the C2 (calibration)
94     + threshold-in-loop task list.
95
96     ### Eval-honesty helpers (course corrections C1?C4) ? additive, default-OFF
*(2026-06-14)*
97     **Why.** A multi-agent literature review (folded into
98     [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) §13) confirmed the
framework's *skeleton* is
99     standard prior art (Temporal Action Localization/Segmentation + late/score-level fusion
? cite, don't claim),
100    and flagged that our **measurement** is walking into known small-N traps that would make
"promote only on
101    measured lift" promote noise or hide a real win. Four corrections, shipped as **pure,
additive helpers** (no
102    existing `compare`/LOVO/`metrics.score` behaviour changed ? callers opt in):
103    - **C1 ? never a single mean.** `backend/eval/eval_stats.py`:
`mean_with_ci`/`bootstrap_ci` (deterministic,
104    seeded), `paired_sign_test` (the exact distribution-free "won on enough *videos*?"
test the corpus already
105    needs), `paired_delta_ci` (B?A with a CI that must clear 0). At n?6 a lone LOVO mean
is not trustworthy.
106    - **C2 ? calibrate cue confidence.** `backend/eval/score_calibration.py`: Platt +
isotonic (PAVA) calibrators
107    + `brier_score`/`expected_calibration_error`, so a CNN softmax, an audio onset, and an
integer net-crossing
108    count are comparable before the arbiter and the over-confident cue can't swamp the
rest.
109    - **C3 ? honest stacking.** `eval_stats.out_of_fold_predictions`/`grouped_folds` (leave-
one-group-out by
110    match) so a *learned* producer's outputs used as arbiter features are out-of-fold (no
leakage).

```

```

111     - C4 ? over-segmentation-visible metrics.** `backend/eval/segmentation_metrics.py`:
`f1_at_overlaps`
112     (F1@{10,25,50}), `average_precision`/`mean_average_precision` (mAP@tIoU),
`segment_count_ratio`. tIoU alone
113     is blind to a rally split in 3 or two rallies merged ? exactly what the reconciler
exists to fix.
114
115     Status. Helpers + 24 unit tests landed (pure stdlib, run in the numpy/GPU-free
lanes); ruff + mypy clean.
116     Not yet wired into the default harness flow (wiring `compare`/`run_regression.py` to
report CIs + F1@k is
117     the follow-on, owner-gated). The discipline they encode is the methodology for every
future cue measurement ?
118     see §13/C1?C4 for the citations and the remaining (Tier-2/3) corrections.
119
120     Multi-signal fusion P2: the FusionSegmenter *(in progress, 2026-06-14)*
121     Hypothesis. Adjudicating shuttle-seeded windows with the net-crossing Gate A
(`min_crossings?2`)
122     and a person-occupancy Gate B raises precision on the held-shuttle / warm-up /
one-sided-feed /
123     spectator false positives ? without hurting recall, because the shuttle stays the
primary,
124     camera-invariant seed.
125
126     Scope ? and an honesty boundary from the archive (read this).** The
127     [2026-06-multivideo-lovo](archives/quality-iterations/2026-06-multivideo-
lovo/README.md) iteration
128     already smoke-tested court-region detection as a fix for the multi-court contrast
confound** and
129     found it low-headroom: only 12?14% of dead-time shuttle detections fall outside
the foreground
130     court, so a perfect court gate lifts contrast to only ~1.6?2.2× (below the 3.6× "works"
threshold) ? the
131     multi-court problem is temporal, and the learned model, not a court mask, is its
answer. So Gate B
132     here is NOT sold as the multi-court fix. It targets a different, previously-
untested failure mode
133     (held-shuttle/warm-up/one-sided/spectator FPs). Two distinct levers; don't conflate
them.
134
135     Data support for Gate A. The same archive shows real rallies on the multi-court
matches contain
136     22 net-crossings 97% of the time (median 5)** while a single service feed crosses
~once ? min_crossings?2
137     is a principled, data-backed discriminator (it's the held-shuttle owner-bug killer).
138
139     Shipping shape. A NEW fusion segmenter, default OFF (default_segmenter
unchanged); Gate A
140     wired into the live trajectory path behind indexing.min_crossings (default 0 = OFF ?
wasb/tracknet
141     byte-identical today); the court mask is an optional config polygon (omit ? Gate B
is
142     player-count-only; supply ? masked). Person features (players_in_court,
two_sided_motion,
143     mean_player_motion, in_court_ratio, shuttle_player_colocation) persisted into candidate
stats for P3's
144     learned fusion. Reuses PersonDetector (P1) + rally_gate (already existed) ? no
reinvention.
145
146     Status (2026-06-14): SHIPPED, default-OFF. FusionSegmenter
(backend/pipeline/segmenters/fusion_hybrid.py,
147     registry key fusion_hybrid) + the pure feature math
(backend/pipeline/detectors/fusion_features.py) + two
148     identity hooks on the trajectory engine (enrich_candidate_stats,
select_for_handoff) + PersonDetector.

```

```

149     detections_per_frame` + the `FusionConfig` block. Built design-first (a
`fusion-p2-design` fan-out workflow), then
150     **adversarially reviewed by a 31-agent workflow (25 findings ? 19 confirmed)**: the
headline confirmation is
151     **NO-REGRESSION verified** ? wasb/tracknet twins persist byte-identical stats and hand
off all windows (pinned by
152     the extended equivalence test); default `FusionSegmenter` == its wasb substrate. The
confirmed low/medium findings
153     were fixed (frame-dim guards, `round()` frame indices, the per-window axis-estimate
redundancy ? a reusable
154     `rally_gate` estimate, honest torch-import docstrings, and 6 new tests closing the
person-detector / Gate-B /
155     substrate / colocation-halo gaps). Suite 441?467 green**; ruff + mypy clean.
156     **Two findings deferred with intent** (reviewer-agreed): person features are computed
eagerly even for
157     Gate-A-pruned windows (a P3 redesign, would change the persisted-superset semantics);
and `fusion.velocity_threshold`
158     must be set equal to the substrate's to A/B against a tuned wasb run (documented in
`FusionConfig`).
159     **Lift still to be measured** offline on the 6 golden videos via
`experiment.compare(CONTROL, fusion)`; the
160     held-shuttle-drop + person-feature values are GPU/owner-side. *(Measured numbers
appended here when run ? no claims before.)*
161
162     ### Person cue extracted to a reusable producer (fusion P1) ?
[#115](https://github.com/avidullu/badminton-highlight-indexer/pull/115), 2026-06-13
163     Lifted the torchvision detector + ByteTrack + velocity windowing out of `yolo_hybrid`
into
164     `backend/pipeline/detectors/person_detector.py::PersonDetector`. **No behavior change**
(proven by the
165     unchanged `process_video` output tests); `yolo_hybrid` now orchestrates only. **Why it's
a quality move:**
166     it makes "add the person cue to fusion" an *add-a-producer*, not a segmenter rewrite.
Suite 440?441 green.
167
168     ### Experiment harness: A/B without rollback risk ?
[#114](https://github.com/avidullu/badminton-highlight-indexer/pull/114), 2026-06-13
169     Shipped `backend/eval/experiment.py` ? the thin variant/experiment layer (the audit's
"-80% already
170     exists" held; pure glue over
`metrics`/`calibration`/`classifier`/`training`/`regression.gt_hash`). Gives
171     `Variant` (+ pinned `CONTROL`), `compare()` (labeled A/B, LOVO-honest),
**compare_unlabeled()`** (the SxS
172     read on NEW unlabeled footage ? agreed / A-only / B-only), and a JSON experiment
leaderboard. **This is
173     the instrument every iteration is measured on.** 19 tests incl. the no-regression
invariant (all cues OFF
174     ? byte-identical to control). Suite 421?440 green.
175
176     ---
177
178     ## Archived predecessors (the historical trail)
179     Completed iterations live under [archives/quality-iterations/](archives/quality-
iterations/README.md):
180     - **2026-06-initial/** ? seed of the quality phase.
181     - **2026-06-human-gt-tuning/** ? first eval vs **human** GT on a partial-label video;
partial-label
182     harness + labeled-window methodology; baseline 0.650 ? tuned 0.742 (*fit-on-self*);
found+fixed a
183     negative-start ingestion bug (collector PR #13).
184     - **2026-06-multivideo-lovo/** ? the **n=6 corpus + honest LOVO**; established the
heuristic is
185     **footage-fragile** (multi-court Kushagra 0.16 / gBaaddy 0.28), the **"contrast"
metric** predicts
186     heuristic F1 monotonically, and a **learned per-frame prototype overtook the heuristic

```

```

at n=3 and holds
187     ~+0.10** (n=6: learned 0.583 vs heuristic 0.480), with **AUC 0.83?0.92 decoupled from
contrast**. This
188     is the data-backed green light for the owned model.
189
190     ---
191
192     ## Standing lessons (carry-forward)
193     - **Persist detection once, re-score variants forever.** Expensive/risky = per-frame
detection (GPU/WSL);
194     cheap/swappable = the decision. Most experiments never enter the served pipeline.
195     - **LOVO or it didn't happen.** Score on a held-out *video*, never a held-out window
from the same video.
196     Group by **match**, not file. Sign test needs n ? ~n=10/9-wins for p<0.05.
197     - **The multi-court confound is TEMPORAL, not spatial/loudness** ? three heuristics were
tested and failed
198     for the same reason (foreground spatial gate **?0.04**; audio energy **AUC ~0.6**;
court-region
199     separability **12?14% out-of-region**): the dead-time confusion is rally-vs-non-rally
activity in the
200     *same place and frequency band*. Only the learned model's temporal/structural features
crack it (AUC
201     0.83 on Kushagra/gBaaddy). **Don't re-derive these dead ends.**
202     - **Court-masking ? multi-court fix** (see above) ? it's a **precision** tool for off-
court persons and
203     one-sided feeds, a *different* lever; measure its lift, don't assume it.
204     - **Net-crossings are a strong rally signal** ? 97% of real rallies cross ?2x (median
5); single feeds ~1x.
205     Gate A leans on this.
206     - **Small-N feature discipline.** Few, scale/translation-invariant features
(counts/ratios, not raw
207     coords) + a logistic regression, not a net ? raw coords memorize this court/camera.
208     - **Calibration, not representation, is the small-N wall** ? learned probabilities
separate (high AUC) but
209     the operating point doesn't transfer with few training videos. More distinct matches
is the highest-ROI
210     unblock.
211     - **Gate A already existed, just wasn't wired to production** ? the cheapest precision
win is plumbing, not
212     new capability. Always check what the eval drivers already do that the live path
doesn't.
213

```

21. assistant (2026-06-14T09:50:18.460Z)

I now have full context. Let me look at the first-A/B archive for the exact Gate-A numbers and over-seg ratio, and check the `distill\_local` feature names. Then I'll build and run the leave-one-signal-out ablation.

22. user (2026-06-14T09:50:23.158Z)

```

README.md
---
FEATURE_NAMES: ['duration', 'peak_velocity', 'mean_velocity', 'active_ratio', 'net_crossings']
---audio---
AUDIO: ['audio_hit_count', 'audio_hit_rate', 'audio_hit_strength_mean']
---person---
PERSON: ['players_in_court', 'two_sided_motion', 'mean_player_motion', 'in_court_ratio',
'shuttle_player_colocation']

```

23. user (2026-06-14T09:50:31.284Z)

```

1     # First end-to-end A/B experiment on the golden corpus (n=6)
2

```

```

3      **STATUS: COMPLETED (2026-06-14).** The first time the experiment harness
4      ([EXPERIMENT_HARNESS.md](../.../EXPERIMENT_HARNESS.md)) + the measurement-honesty
helpers
5      (course corrections C1?C4,
[ANALYZERS_AND_RECONCILERS.md](../.../ANALYZERS_AND_RECONCILERS.md) §13)
6      were run end-to-end on real golden data ? a posterity record of the framework's first
real use.
7
8      **Trigger.** After landing the signal-fusion framework doc + §13 literature review and
the additive
9      C1?C4 helpers (`backend/eval/eval_stats.py`, `score_calibration.py`,
`segmentation_metrics.py`, PR #124),
10     the question was: *does the framework actually hold up on a golden video?* This is the
validation run.
11
12     ## Method (100% offline ? the "persist detection once, re-score forever" path)
13     - **No GPU.** Re-scored the **persisted WASB shuttle trajectories** of the 6 golden
matches
14     (`output/human_lovo_manifest.json`); detection was never re-run.
15     - **One consistent candidate set per video** via `distill_local.build_candidates`
(recall-oriented
16     proposal windowing `velocity_thresh=0.005, merge_gap=1.0, min_dur=0.5`) ? per-window
shuttle features
17     `[duration, peak_velocity, mean_velocity, active_ratio, net_crossings]`; golden labels
via
18     `calibrate_wasb.load_golden_gts`; per-candidate y via the evaluator's own IoU matcher.
19     - **Three variants re-scored on that same set** (`backend/eval/experiment.py`):
20     - **CONTROL** ? merge all candidates (no filter).
21     - **Gate-A** ? keep candidates with `net_crossings` ? 2, then merge (the held-shuttle
/ service-feed
22     structural FP killer; `rally_gate`).
23     - **learned** ? a window-level logistic regression over the 5 shuttle features,
**leave-one-VIDEO-out**
24     (train on the other 5 matches, predict held-out), prob ? 0.5, merge.
25     - **Graded with:** temporal-IoU F1 (per video), **C1** honest aggregation
(`mean_with_ci`,
26     `paired_delta_ci`, exact `paired_sign_test`), **C4** segmentation metrics
(`f1_at_overlaps`,
27     `segment_count_ratio`), and a **C3** out-of-fold learned mAP@tIoU.
28     - IoU 0.5. Repro: `scratch/ab_n6_lovo.py` (machine-side; reads the gitignored manifest +
the external
29     golden trajectories/labels); raw result `output/ab_experiment_n6.json`.
30
31     ## Corpus
32     | video | candidates | golden rallies | proposal recall ceiling |
33     |---|---|---|---|
34     | adarsh_avi_singles | 172 | 96 | 0.833 |
35     | kushagra_singles (multi-court) | 20 | 32 | **0.125** |
36     | largetest_doubles | 86 | 49 | 0.816 |
37     | gbaaddy (multi-court) | 82 | 48 | **0.312** |
38     | testlarge_short | 14 | 6 | 0.500 |
39     | mahadevpura_singles | 57 | 31 | 0.645 |
40
41     The multi-court matches (kushagra, gbaaddy) have **low proposal-recall ceilings** ? the
decision stage
42     can never recover rallies the windowing missed (the proposal-recall ceiling pitfall,
§13). Their absolute
43     F1 is windowing-capped, not a decision failure.
44
45     ## Results
46
47     **Per-video held-out F1 (IoU 0.5):**
48     | video | CONTROL | Gate-A | learned (LOVO) |
49     |---|---|---|---|
50     | adarsh_avi_singles | 0.597 | **0.722** | 0.613 |

```

```

51 | kushagra_singles | 0.154 | **0.163** | 0.163 |
52 | largetest_doubles | 0.593 | **0.667** | 0.535 |
53 | gbaaddy | 0.231 | **0.274** | 0.222 |
54 | testlarge_short | 0.300 | **0.375** | **0.000** |
55 | mahadevpura_singles | 0.455 | **0.571** | 0.308 |
56 | **mean** | **0.388** | **0.462** | **0.307** |
57
58 **A/B 1 ? CONTROL vs Gate-A (static).** ?F1 **+0.074**, 95% CI **[+0.042, +0.104]**
(excludes 0),
59 sign test **6 wins / 0 losses, p = 0.031**. ? **A real, statistically-supported win** ?
Gate-A improves F1
60 on *every* golden match.
61
62 **A/B 2 ? CONTROL vs window-level learned logistic (LOVO).** ?F1 **?0.081**, CI
**[?0.173, ?0.002]**,
63 sign test **2 wins / 4 losses, p = 0.688**. ? **Does not clear the bar** ? this variant
is *worse* than
64 CONTROL on this setup (and the mean is dragged by two big per-video losses; the sign
test is inconclusive ?
65 the two honest views disagreeing is itself a reason to report both).
66
67 **C4 ? over-segmentation-aware metrics (mean across videos):**
68 | variant | F1@10 | F1@25 | F1@50 | segment-count ratio |
69 |---|---|---|---|---|
70 | CONTROL | 0.645 | 0.603 | 0.388 | 1.68 |
71 | **Gate-A** | **0.748** | **0.698** | **0.462** | **1.01** |
72 | learned | 0.564 | 0.509 | 0.307 | 0.89 |
73
74 Gate-A drives the **segment-count ratio 1.68 ? 1.01** ? it nearly eliminates the over-
segmentation that
75 plain temporal-IoU barely reflected. The learned variant's 0.89 is *under*-segmentation
(it drops real
76 rallies). **C3 out-of-fold learned mAP@tIoU[.3?.7] = 0.245** (weak ranking ?
corroborates the loss).
77
78 ## Wins / losses / lessons
79 - ? **Gate-A (net-crossings ? 2) is a genuine win** ? +0.074 mean F1, 6/6 matches, p =
0.031, CI clears 0,
80 over-segmentation 1.68?1.01. Promote-worthy by the honest gate. (Confirms the held-
shuttle owner-bug
81 killer on the whole corpus, offline, no GPU.)
82 - ? **The window-level 5-feature logistic (threshold 0.5) is not a win here** ? ?0.081,
mixed (2W/4L).
83 `testlarge_short ? 0.000` (the model predicted *zero* rallies) is a
**threshold/calibration artifact**, a
84 textbook motivation for **C2 (calibrate cue confidence)** and tuning the threshold
*inside* the LOVO loop.
85 - ?? **Do not conflate this with the documented learned winner.** The
86 [2026-06-multivideo-lovo](../2026-06-multivideo-lovo/README.md) result (learned
0.583 vs heuristic
87 0.480) is a **per-frame prototype** (`rally_seq_proto.py`), a *different, stronger*
model than this
88 window-level logistic. This experiment compared a *weaker* learned path against a
*strong* CONTROL
89 (merge-all of a recall-oriented proposal) ? a different matchup. The takeaway is
methodological, not
90 "learned fusion fails."
91 - ? **The methodology itself worked.** C1 gave a confident **YES** on Gate-A *and* a
clear **NO** on the
92 logistic ? preventing a wrong promotion in **both** directions. A bare-mean report
would have hidden
93 Gate-A's significance and might have over-read the logistic's mean. C4 surfaced the
over-segmentation
94 story tIoU buried. This is exactly the framework holding up.
95 - ? **Multi-court is windowing-capped** ? kushagra/gbaaddy absolute F1 is bounded by the

```

```

0.125/0.312
96     proposal-recall ceiling; the decision layer can't fix what the proposal missed
(measure AR@AN separately).
97
98     ## Significance
99     First real A/B on the golden corpus ? the harness + C1?C4 produced an **honest,
promotable finding
100     (Gate-A)** and an **honest rejection (this logistic)**, entirely offline. This is the
empirical backbone the
101     paper aim ([`paper-coauthor-aim`]) needs ? honest measured results only. Gate-A becomes
a candidate for the
102     no-regression promotion gate; the logistic's failure mode is the concrete task list for
C2 + threshold-in-loop.
103
104     **Links:** [EXPERIMENT_HARNESS.md](../../EXPERIMENT_HARNESS.md) ·
105     [ANALYZERS_AND_RECONCILERS.md](../../ANALYZERS_AND_RECONCILERS.md) §13 (C1?C4) ·
106     [QUALITY_ITERATIONS.md](../../QUALITY_ITERATIONS.md) · predecessor
107     [`2026-06-multivideo-lovo`](../2026-06-multivideo-lovo/README.md).
108

```

24. assistant (2026-06-14T09:50:38.181Z)

I have everything I need to design the ablation. Key understanding:

- **Feature columns** (consumed by the learned logistic): shuttle (5: duration, peak_velocity, mean_velocity, active_ratio, net_crossings), person (5), audio (3) = 13 total.
- **Gates** (decision masks, not learned columns): Gate-A = `min\_crossings` on `net\_crossings`; Gate-B = `min\_players` on `players\_in\_court`.

The owner wants a **leave-one-signal-out** ablation. For the learned-fusion layer, the natural "full set" = all 13 feature columns; leave-one-group-out removes each signal group. For the gates, Gate-A is a re-scoring mask, so I'll measure Gate-A ON vs OFF separately (its known result + reproduction + over-seg ratio). Let me write the ablation script in scratch and run it.

Let me mark this as a new chapter and write the ablation runner.

25. user (2026-06-14T09:50:39.400Z)

Chapter marked: "Gate-A/B leave-one-signal-out ablation". Continue your current work.

26. assistant (2026-06-14T09:51:39.402Z)

""Leave-one-signal-out ablation for the Gate-A/B foundational decision (CPU, offline).

Reuses ONLY the shipped harness (experiment.compare + eval_stats) over the 6 persisted golden feature JSONs (output/*_golden_features.json). No new scoring math, no GPU, no Gemini.

Two kinds of question, two mechanisms (mirrors the repo's own distinction):

1. **LEARNED-FUSION marginal contribution (leave-one-signal-out):** the "full" learned variant uses ALL 13 feature columns (shuttle 5 + person 5 + audio 3). For each signal GROUP (and each individual shuttle feature where sensible) we drop those columns and measure held-out ?F1 = full - leave_one_out, with bootstrap 95% CI + paired sign test. ?F1 > 0 means the signal carries marginal information the learned scorer used.

NOTE the sign convention: we report delta = FULL - DROPPED, so a POSITIVE delta means the signal HELPS (dropping it hurts). experiment.compare reports B - A, so we set variant_a = leave_one_out (dropped), variant_b = full -> B-A = full - dropped. Good.
2. **GATE marginal contribution (decision masks, NOT learned columns):** Gate-A (min_crossings>=2) and Gate-B (min_players>=N) are re-scoring masks over the recall-oriented CONTROL proposal set. Measured as static-variant A/B: CONTROL vs CONTROL+gate. Reports F1, ?F1 (CI+sign), and the segment-count / over-seg ratio (C4) that tIoU hides.

All numbers printed are ACTUAL measured values. Writes a JSON blob to scratch/ for the doc.

"""

```
from __future__ import annotations
```

```
import json
```

```
import os
```

```
from typing import Any, Dict, List, Sequence
```

```
from backend.eval import experiment
```

```
from backend.eval.eval_stats import paired_delta_ci
```

```
from backend.eval import segmentation_metrics as segm
```

```
from backend.eval.metrics import score
```

```
from backend.pipeline.detectors import audio_features as af
```

```
from backend.pipeline.detectors import fusion_features as ff
```

```
from backend.eval import distill_local
```

```
SHUTTLE = list(distill_local.FEATURE_NAMES)          # 5
```

```
PERSON = list(ff.PERSON_FEATURE_NAMES)               # 5
```

```
AUDIO = list(af.AUDIO_FEATURE_NAMES)                 # 3
```

```
ALL_FEATURES = SHUTTLE + PERSON + AUDIO              # 13
```

```
FEATURES_DIR = "output"
```

```
IOU = 0.5
```

```
THRESHOLD = 0.5
```

```
def load_videos() -> List[experiment.VideoCandidates]:
```

```
    import glob
```

```
    videos: List[experiment.VideoCandidates] = []
```

```
    for path in sorted(glob.glob(os.path.join(FEATURES_DIR, "*_golden_features.json"))):
```

```
        with open(path, encoding="utf-8") as f:
```

```
            d = json.load(f)
```

```
            cands = d["candidates"]
```

```
            intervals = [(float(c["start"]), float(c["end"])) for c in cands]
```

```
            feats: Dict[str, List[float]] = {}
```

```
            for fn in SHUTTLE:
```

```
                feats[fn] = [float(c["shuttle"][fn]) for c in cands]
```

```
            for fn in PERSON:
```

```
                feats[fn] = [float(c["person"][fn]) for c in cands]
```

```
            for fn in AUDIO:
```

```
                feats[fn] = [float(c["audio"][fn]) for c in cands]
```

```
            gts = [(float(a), float(b)) for a, b in d["gts"]]
```

```
            videos.append(experiment.VideoCandidates(name=d["name"], intervals=intervals,
                                                         features=feats, gts=gts))
```

```
    return videos
```

```
def paired_from_compare(res: Dict[str, Any], metric: str = "f1") -> Dict[str, Any]:
```

```
    """Pair B-A per-video held-out scores -> mean delta + CI + sign test."""
```

```
    by_a = {p["video"]: p[metric] for p in res["variant_a"]["per_video"]}
```

```
    by_b = {p["video"]: p[metric] for p in res["variant_b"]["per_video"]}
```

```
    vids = [p["video"] for p in res["variant_a"]["per_video"]]
```

```
    return paired_delta_ci([by_a[v] for v in vids], [by_b[v] for v in vids])
```

```
def learned_lovo(videos, feature_names) -> Dict[str, Any]:
```

```
    """Run a learned LOVO variant over `feature_names`; return aggregate + per-video f1."""
```

```
    v = experiment.Variant(name="v", feature_names=list(feature_names), threshold=THRESHOLD)
```

```
    ev = experiment.evaluate_variant(videos, v, iou=IOU, honest=True)
```

```
    return ev
```

```
def leave_one_out_row(videos, full_features, drop, label) -> Dict[str, Any]:
```

```
    """full (all features) as B; leave-one-out (full minus `drop`) as A.
```

```
    delta = B - A = full - dropped -> positive means the dropped signal HELPS."""
```

```

kept = [f for f in full_features if f not in set(drop)]
a = experiment.Variant(name=f"drop_{label}", feature_names=kept, threshold=THRESHOLD)
b = experiment.Variant(name="full", feature_names=list(full_features), threshold=THRESHOLD)
res = experiment.compare(videos, a, b, iou=IOU, honest=True) # B-A = full - dropped
pd = paired_from_compare(res)
st = pd["sign_test"]
return {
    "signal": label,
    "dropped_cols": list(drop),
    "f1_full": res["variant_b"]["aggregate"]["f1"],
    "f1_dropped": res["variant_a"]["aggregate"]["f1"],
    "marginal_dF1": pd["mean_delta"], # full - dropped
    "ci_low": pd["ci_low"], "ci_high": pd["ci_high"],
    "ci_excludes_zero": pd["ci_excludes_zero"],
    "wins": int(st["wins"]), "losses": int(st["losses"]),
    "p_value": st["p_value"],
    "per_video_full": {p["video"]: p["f1"] for p in res["variant_b"]["per_video"]},
    "per_video_dropped": {p["video"]: p["f1"] for p in res["variant_a"]["per_video"]},
}

def static_predict_segcount(variant, videos) -> Dict[str, float]:
    """Mean F1@k + segment_count_ratio for a static variant across videos."""
    overlaps = segm.DEFAULT_OVERLAPS
    f1k = {ov: [] for ov in overlaps}
    ratios = []
    for vc in videos:
        preds = experiment.predict(variant, vc)
        got = segm.f1_at_overlaps(preds, vc.gts or [], overlaps)
        for ov in overlaps:
            f1k[ov].append(got[ov])
        ratios.append(segm.segment_count_ratio(preds, vc.gts or []))
    out = {f"f1@{int(ov*100)}": (sum(v)/len(v) if v else 0.0) for ov, v in f1k.items()}
    out["segment_count_ratio"] = sum(ratios)/len(ratios) if ratios else 0.0
    return out

def gate_ab(videos, variant_a, variant_b, label) -> Dict[str, Any]:
    """Static A/B (no learned filter): CONTROL vs CONTROL+gate. Reports F1, ?F1 (CI+sign),
    and over-seg ratio for both."""
    res = experiment.compare(videos, variant_a, variant_b, iou=IOU, honest=True)
    pd = paired_from_compare(res)
    st = pd["sign_test"]
    seg_a = static_predict_segcount(variant_a, videos)
    seg_b = static_predict_segcount(variant_b, videos)
    # count of predicted segments per variant (for over-seg narrative)
    def total_preds(v):
        return sum(len(experiment.predict(v, vc)) for vc in videos)
    return {
        "label": label,
        "f1_a": res["variant_a"]["aggregate"]["f1"],
        "f1_b": res["variant_b"]["aggregate"]["f1"],
        "p_a": res["variant_a"]["aggregate"]["precision"],
        "p_b": res["variant_b"]["aggregate"]["precision"],
        "r_a": res["variant_a"]["aggregate"]["recall"],
        "r_b": res["variant_b"]["aggregate"]["recall"],
        "dF1": pd["mean_delta"], "ci_low": pd["ci_low"], "ci_high": pd["ci_high"],
        "ci_excludes_zero": pd["ci_excludes_zero"],
        "wins": int(st["wins"]), "losses": int(st["losses"]), "p_value": st["p_value"],
        "segratio_a": seg_a["segment_count_ratio"], "segratio_b": seg_b["segment_count_ratio"],
        "n_pred_a": total_preds(variant_a), "n_pred_b": total_preds(variant_b),
        "per_video_dF1": {p["video"]: p["f1"] for p in res["per_video_delta"]},
        "f1k_a": seg_a, "f1k_b": seg_b,
    }

```

```

def main() -> None:
    videos = load_videos()
    print(f"Loaded {len(videos)} golden videos: {[v.name for v in videos]}")
    result: Dict[str, Any] = {"n_videos": len(videos), "iou": IOU, "threshold": THRESHOLD}

    # ---- Full learned variant (all 13 cols) reference ----
    full_ev = learned_lovo(videos, ALL_FEATURES)
    result["full_learned"] = {
        "feature_names": ALL_FEATURES,
        "aggregate": full_ev["aggregate"],
        "per_video": {p["video"]: p["f1"] for p in full_ev["per_video"]},
    }
    print("\n=== FULL learned LOVO (all 13 feature columns) ===")
    print(f" F1={full_ev['aggregate']['f1']:.3f} P={full_ev['aggregate']['precision']:.3f} "
          f"R={full_ev['aggregate']['recall']:.3f}")

    # ---- Leave-one-SIGNAL-GROUP-out ----
    print("\n=== Leave-one-signal-out marginal contribution (delta = full - dropped) ===")
    groups = [
        ("person", PERSON),
        ("audio", AUDIO),
        ("net_crossings(shuttle-as-feature)", ["net_crossings"]),
        ("shuttle_velocity_block", ["peak_velocity", "mean_velocity", "active_ratio"]),
        ("duration", ["duration"]),
    ]
    rows = []
    for label, drop in groups:
        row = leave_one_out_row(videos, ALL_FEATURES, drop, label)
        rows.append(row)
        print(f" drop {label:<36s} ?F1(full-drop)={row['marginal_dF1']:+.3f} "
              f"CI[{row['ci_low']:+.3f},{row['ci_high']:+.3f}] "
              f"{row['wins']}W/{row['losses']}L p={row['p_value']:.3f} "
              f"clears0={row['ci_excludes_zero']}")
    result["leave_one_out"] = rows

    # ---- Also: incremental ADD-each-group-onto-shuttle (matches P3/P4 framing) ----
    print("\n=== Incremental add-onto-shuttle (B - A) ===")
    add_rows = []
    incr = [
        ("shuttle->+person", SHUTTLE, SHUTTLE + PERSON),
        ("shuttle->+audio", SHUTTLE, SHUTTLE + AUDIO),
        ("shuttle+person->+audio", SHUTTLE + PERSON, SHUTTLE + PERSON + AUDIO),
        ("shuttle->+person+audio", SHUTTLE, ALL_FEATURES),
    ]
    for label, a_f, b_f in incr:
        a = experiment.Variant(name="a", feature_names=a_f, threshold=THRESHOLD)
        b = experiment.Variant(name="b", feature_names=b_f, threshold=THRESHOLD)
        res = experiment.compare(videos, a, b, iou=IOU, honest=True)
        pd = paired_from_compare(res)
        st = pd["sign_test"]
        r = {"label": label, "f1_a": res["variant_a"]["aggregate"]["f1"],
            "f1_b": res["variant_b"]["aggregate"]["f1"], "dF1": pd["mean_delta"],
            "ci_low": pd["ci_low"], "ci_high": pd["ci_high"],
            "ci_excludes_zero": pd["ci_excludes_zero"],
            "wins": int(st["wins"]), "losses": int(st["losses"]), "p_value": st["p_value"]}
        add_rows.append(r)
        print(f" {label:<28s} F1 {r['f1_a']:.3f}->{r['f1_b']:.3f} ?F1={r['dF1']:+.3f} "
              f"CI[{r['ci_low']:+.3f},{r['ci_high']:+.3f}] {r['wins']}W/{r['losses']}L "
              f"p={r['p_value']:.3f} clears0={r['ci_excludes_zero']}")
    result["incremental_add"] = add_rows

    # ---- GATE-A ON vs OFF (static, the foundational decision) ----
    print("\n=== Gate-A ON vs OFF (static CONTROL vs net_crossings>=2) ===")
    control = experiment.Variant(name="control") # merge-all, no gate
    gate_a = experiment.Variant(name="gate_a", min_crossings=2)

```

```

ga = gate_ab(videos, control, gate_a, "Gate-A (min_crossings>=2)")
result["gate_a"] = ga
print(f"  CONTROL F1={ga['f1_a']:.3f} (P={ga['p_a']:.3f} R={ga['r_a']:.3f})  "
      f"segratio={ga['segratio_a']:.2f}  n_pred={ga['n_pred_a']}")
print(f"  Gate-A F1={ga['f1_b']:.3f} (P={ga['p_b']:.3f} R={ga['r_b']:.3f})  "
      f"segratio={ga['segratio_b']:.2f}  n_pred={ga['n_pred_b']}")
print(f"  ?F1={ga['dF1']:.3f} CI[{ga['ci_low']:.3f},{ga['ci_high']:.3f}]  "
      f"{ga['wins']}W/{ga['losses']}L p={ga['p_value']:.3f}  "
      f"clears0={ga['ci_excludes_zero']}")

# ---- GATE-B ON vs OFF (static; min_players sweep). No court polygons -> caveat. ----
print("\n=== Gate-B ON vs OFF (static CONTROL vs players_in_court>=N) ===")
gb_rows = []
for n in (1, 2, 3, 4):
    gb = experiment.Variant(name=f"gate_b_{n}", min_players=n)
    r = gate_ab(videos, control, gb, f"Gate-B (min_players>={n})")
    gb_rows.append(r)
    print(f"  N>={n}: CONTROL F1={r['f1_a']:.3f} -> Gate-B F1={r['f1_b']:.3f}  "
          f"?F1={r['dF1']:.3f} CI[{r['ci_low']:.3f},{r['ci_high']:.3f}]  "
          f"{r['wins']}W/{r['losses']}L p={r['p_value']:.3f} clears0={r['ci_excludes_zero']}")
    result["gate_b"] = gb_rows

# ---- Gate-A AND Gate-B stacked (does B add on top of A?) ----
print("\n=== Gate-A + Gate-B stacked vs Gate-A alone ===")
gab_rows = []
for n in (1, 2):
    a = experiment.Variant(name="gate_a", min_crossings=2)
    b = experiment.Variant(name=f"gate_ab_{n}", min_crossings=2, min_players=n)
    res = experiment.compare(videos, a, b, iou=IOU, honest=True)
    pd = paired_from_compare(res)
    st = pd["sign_test"]
    r = {"n": n, "f1_a": res["variant_a"]["aggregate"]["f1"],
         "f1_b": res["variant_b"]["aggregate"]["f1"], "dF1": pd["mean_delta"],
         "ci_low": pd["ci_low"], "ci_high": pd["ci_high"],
         "ci_excludes_zero": pd["ci_excludes_zero"],
         "wins": int(st["wins"]), "losses": int(st["losses"]), "p_value": st["p_value"]}
    gab_rows.append(r)
    print(f"  GateA + B(N>={n}) vs GateA: F1 {r['f1_a']:.3f}->{r['f1_b']:.3f} "
          f"?F1={r['dF1']:.3f}  "
          f"CI[{r['ci_low']:.3f},{r['ci_high']:.3f}] {r['wins']}W/{r['losses']}L "
          f"p={r['p_value']:.3f}")
    result["gate_ab_stacked"] = gab_rows

out = os.path.join("scratch", "gate_ab_ablation_result.json")
with open(out, "w", encoding="utf-8") as f:
    json.dump(result, f, indent=2)
print(f"\nWrote {out}")

if __name__ == "__main__":
    main()

```

27. user (2026-06-14T09:51:39.420Z)

File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\gate_ab_ablation.py (file state is current in your context ? no need to Read it back)

28. user (2026-06-14T09:51:42.979Z)

Exit code 1

Traceback (most recent call last):

File "C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\gate_ab_ablation.py", line

```
31, in <module>
    from backend.eval import experiment
ModuleNotFoundError: No module named 'backend'
```

29. assistant (2026-06-14T09:51:45.365Z)

Need to run from the repo root with proper module path. The script imports work as a module; let me run with `PYTHONPATH` set to the repo root.

30. user (2026-06-14T09:51:46.640Z)

```
Exit code 1
/usr/bin/bash: line 1: :PYTHONPATH: command not found
Traceback (most recent call last):
  File "C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\gate_ab_ablation.py", line
31, in <module>
    from backend.eval import experiment
ModuleNotFoundError: No module named 'backend'
```

31. assistant (2026-06-14T09:51:49.294Z)

The Bash tool is POSIX (Git bash), not PowerShell. Let me use the PowerShell tool for env var handling, or set it inline.

32. user (2026-06-14T09:51:57.382Z)

Loaded 6 golden videos: ['adarsh_avi_singles', 'gbaaddy', 'kushagra_singles', 'largetest_doubles', 'mahadevpura_singles', 'testlarge_short']

```
=== FULL learned LOVO (all 13 feature columns) ===
F1=0.366 P=0.453 R=0.341
```

```
=== Leave-one-signal-out marginal contribution (delta = full - dropped) ===
drop person                ?F1(full-drop)=-0.009 CI[-0.074,+0.055] 2W/4L
p=0.688 clears0=False
drop audio                  ?F1(full-drop)=+0.079 CI[-0.054,+0.203] 4W/2L
p=0.688 clears0=False
drop net_crossings(shuttle-as-feature) ?F1(full-drop)=-0.014 CI[-0.083,+0.027] 4W/1L
p=0.375 clears0=False
drop shuttle_velocity_block ?F1(full-drop)=+0.002 CI[-0.010,+0.018] 2W/3L
p=1.000 clears0=False
drop duration               ?F1(full-drop)=-0.001 CI[-0.004,+0.000] 0W/2L
p=0.500 clears0=False
```

```
=== Incremental add-onto-shuttle (B - A) ===
shuttle->+person            F1 0.307->0.286 ?F1=-0.021 CI[-0.165,+0.162] 2W/4L p=0.688
clears0=False
shuttle->+audio              F1 0.307->0.375 ?F1=+0.068 CI[-0.002,+0.140] 4W/2L p=0.688
clears0=False
shuttle+person->+audio        F1 0.286->0.366 ?F1=+0.079 CI[-0.054,+0.203] 4W/2L p=0.688
clears0=False
shuttle->+person+audio        F1 0.307->0.366 ?F1=+0.059 CI[-0.039,+0.156] 4W/2L p=0.688
clears0=False
```

```
=== Gate-A ON vs OFF (static CONTROL vs net_crossings>=2) ===
CONTROL F1=0.388 (P=0.313 R=0.539) segratio=1.68 n_pred=431
Gate-A F1=0.462 (P=0.477 R=0.468) segratio=1.01 n_pred=235
?F1=+0.074 CI[+0.042,+0.105] 6W/0L p=0.031 clears0=True
```

```
=== Gate-B ON vs OFF (static CONTROL vs players_in_court>=N) ===
N>=1: CONTROL F1=0.388 -> Gate-B F1=0.388 ?F1=+0.000 CI[+0.000,+0.000] 0W/0L p=1.000
clears0=False n_pred 431->431
N>=2: CONTROL F1=0.388 -> Gate-B F1=0.389 ?F1=+0.000 CI[+0.000,+0.001] 1W/0L p=1.000
clears0=False n_pred 431->430
```

```
N>=3: CONTROL F1=0.388 -> Gate-B F1=0.307 ?F1=-0.081 CI[-0.244,+0.000] 0W/1L p=1.000
clears0=False n_pred 431->274
N>=4: CONTROL F1=0.388 -> Gate-B F1=0.283 ?F1=-0.105 CI[-0.305,+0.002] 1W/2L p=1.000
clears0=False n_pred 431->231
```

=== Gate-A + Gate-B stacked vs Gate-A alone ===

```
GateA + B(N>=1) vs GateA: F1 0.462->0.462 ?F1=+0.000 CI[+0.000,+0.000] 0W/0L p=1.000
GateA + B(N>=2) vs GateA: F1 0.462->0.462 ?F1=+0.000 CI[+0.000,+0.000] 0W/0L p=1.000
```

Wrote scratch\gate_ab_ablation_result.json

33. assistant (2026-06-14T09:52:07.651Z)

The ablation ran cleanly with real measured numbers. Key findings:

- **Gate-A ON vs OFF (static, the foundational decision)**:** ?F1 **++0.074****, CI **++[+0.042, +0.105]**** clears 0, ****6W/0L, p=0.031**** ? reproduced exactly. Over-seg ratio ****1.68?1.01****, n_pred 431?235.
- **Gate-B**:** at the production-relevant counts it's a no-op (N?1/N?2: nearly every window already has ?1-2 detected persons because no court polygons ? spectators counted) and at N?3/4 it HURTS (drops real rallies). Confirms the "counts spectators" caveat. Gate-B adds nothing on top of Gate-A.
- **Leave-one-out (learned columns)**:** only ****audio**** carries positive marginal info (+0.079, but CI spans 0); person ? neutral/slightly negative; net_crossings-as-a-column ? neutral (its power lives in the GATE, not the column); velocity/duration ? zero marginal.

One thing to verify: the Gate-A static A/B uses the same recall-oriented CONTROL proposal set as the archive. Let me confirm the CONTROL F1 (0.388) matches the archive and that the gate's `merge\_gap` default matches. It does (0.388 / 0.462 / segratio 1.68?1.01 all match the archive exactly).

Now let me also double-check one subtlety: the leave-one-out "full" learned F1 is 0.366, matching the P4 endpoint. Good ? internally consistent. Let me verify the proposal-recall ceiling numbers per video and pull the per-video Gate-A deltas for the table.

34. user (2026-06-14T09:52:13.364Z)

Exit code 1

python.exe: ScriptBlock should only be specified as a value of the Command parameter.

35. assistant (2026-06-14T09:52:16.596Z)

The multi-line `python -c` with the for-loop f-strings tripped PowerShell parsing. Let me write it to a small scratch file instead.

36. assistant (2026-06-14T09:52:21.104Z)

```
import json
d = json.load(open('scratch/gate_ab_ablation_result.json', encoding='utf-8'))
ga = d['gate_a']
print('Gate-A per-video DeltaF1 (B-A):')
for k, v in ga['per_video_dF1'].items():
    print(f' {k:<24s} {v:+.3f}')
print()
print('CONTROL F1@k:', {k: round(v, 3) for k, v in ga['f1k_a'].items()})
print('Gate-A F1@k:', {k: round(v, 3) for k, v in ga['f1k_b'].items()})
print()
print('Audio leave-out per-video (full-13 vs drop-audio):')
for r in d['leave_one_out']:
    if r['signal'] == 'audio':
        for vid in r['per_video_full']:
            print(f" {vid:<24s} full={r['per_video_full'][vid]:.3f} drop-
audio={r['per_video_dropped'][vid]:.3f}")
```

```

print()
print('Person leave-out per-video (full-13 vs drop-person):')
for r in d['leave_one_out']:
    if r['signal'] == 'person':
        for vid in r['per_video_full']:
            print(f" {vid:<24s} full={r['per_video_full'][vid]:.3f} drop-
person={r['per_video_dropped'][vid]:.3f}")

```

37. user (2026-06-14T09:52:21.118Z)

File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\gate_ab_detail.py (file state is current in your context ? no need to Read it back)

38. user (2026-06-14T09:52:27.634Z)

Gate-A per-video DeltaF1 (B-A):

adarsh_avi_singles	+0.125
gbaaddy	+0.043
kushagra_singles	+0.009
largetest_doubles	+0.074
mahadevpura_singles	+0.117
testlarge_short	+0.075

CONTROL F1@k: {'f1@10': 0.645, 'f1@25': 0.603, 'f1@50': 0.388, 'segment_count_ratio': 1.675}

Gate-A F1@k: {'f1@10': 0.748, 'f1@25': 0.698, 'f1@50': 0.462, 'segment_count_ratio': 1.014}

Audio leave-out per-video (full-13 vs drop-audio):

adarsh_avi_singles	full=0.677 drop-audio=0.645
gbaaddy	full=0.102 drop-audio=0.000
kushagra_singles	full=0.105 drop-audio=0.128
largetest_doubles	full=0.596 drop-audio=0.324
mahadevpura_singles	full=0.492 drop-audio=0.222
testlarge_short	full=0.222 drop-audio=0.400

Person leave-out per-video (full-13 vs drop-person):

adarsh_avi_singles	full=0.677 drop-person=0.568
gbaaddy	full=0.102 drop-person=0.250
kushagra_singles	full=0.105 drop-person=0.150
largetest_doubles	full=0.596 drop-person=0.614
mahadevpura_singles	full=0.492 drop-person=0.525
testlarge_short	full=0.222 drop-person=0.143

39. assistant (2026-06-14T09:52:39.813Z)

All numbers confirmed and internally consistent. Per-video Gate-A is positive on all 6 (6W/0L confirmed at the per-video level). Now I have everything for the comprehensive table. Let me assess the Gemini question before writing the doc.

****Gemini decision reasoning****: The persisted-feature ablation answers the fusion-layer question (which signals carry marginal information, and whether Gate-A as a decision mask is a real win) decisively and at zero cost. A served end-to-end Gemini run would measure something **different** ? final highlight-reel quality with Gate-A flipped on in the live served path. But:

1. The decision under consideration is the ****Gate-A default flip in the fusion/proposal layer****, which is upstream of Gemini and already measured against golden labels (6/6, p=0.031).
2. There are ****no golden labels**** for any new served-pipeline footage, so a Gemini run produces divergence (a_only/b_only), not F1 ? it can't move a numbers-driven decision, only illustrate it.
3. The known GX010125 offline re-score (45?27 rallies, 18 low-crossing windows dropped) ****already corroborates**** the served-side behavior at zero Gemini cost.

So a served Gemini run is ****not warranted**** to make this decision. I'll flag it with expected cost/time for the owner as an optional confidence-builder. Let me write the doc.

Gate-A/B foundational decision ? golden-set numbers + implications

****Owner's framing:**** *NUMBERS INFORM, FOUNDATIONS DECIDE.* Extensibility, scope of exploration, and rigid foundational choices weigh more than raw ?F1 ? but this must be a ****data-INFLUENCED**** choice, so every number below is ****actually measured**** (no fabrication) on the n=6 owned golden corpus, IoU ? 0.5, with ****honest stats** (C1: bootstrap 95% CI + paired sign test; never a bare mean at n=6)**.

****Repro (CPU, seconds, no GPU/Gemini):****
`PYTHONPATH=<repo> python scratch/gate\_ab\_ablation.py` ? `scratch/gate\_ab\_ablation\_result.json`.
Reuses ONLY the shipped harness (`backend/eval/experiment.compare` + `backend/eval/eval\_stats`) over
the 6 persisted `output/\*\_golden\_features.json`. All 6 JSONs already carry the audio columns ?
no
`--augment` step was needed.

****Corpus**** (the 6 labelled matches in `output/human\_lovo\_manifest.json`):

video	candidates	golden rallies	proposal-recall ceiling	note
adarsh_avi_singles	172	96	0.833	clean singles
largetest_doubles	86	49	0.816	doubles
mahadevpura_singles	57	31	0.645	single-court-ish
gbaaddy	82	48	**0.312**	multi-court
testlarge_short	14	6	0.500	short
kushagra_singles	20	32	**0.125**	multi-court

Multi-court matches are ****windowing-capped**** (the decision layer cannot recover rallies the proposal windowing missed) ? their low absolute F1 is a proposal-recall problem, not a gate/fusion failure.

A. The two mechanisms (this is the foundational distinction)

The golden JSONs hold ****two structurally different things****, and the Gate-A/B decision lives in the first one:

- ****GATES = decision masks**** over the recall-oriented CONTROL proposal set. `Gate-A = net\_crossings ? 2`
and `Gate-B = players\_in\_court ? N` are hard keep/drop rules applied **before** (or instead of) any learned scorer. In the harness these are `Variant.min\_crossings` / `Variant.min\_players` (`\_gate\_mask` in `experiment.py`). ****This is what the foundational decision is about.****
- ****FEATURE COLUMNS = learned-fusion inputs.**** shuttle (5: `duration, peak\_velocity, mean\_velocity, active\_ratio, net\_crossings`) + person (5) + audio (3) = ****13 columns**** a LOVO LogisticRegression weights. Here a signal contributes "softly" ? the scorer decides how much to trust it.
Removing one group and re-measuring held-out ?F1 = that group's ****marginal contribution****.

Note `net\_crossings` appears in BOTH: as a ****gate**** (hard `?2`, the Gate-A decision) and as a ****soft feature column****. They are not the same lever ? and the numbers below show the **gate** form is where all the power is.

B. Comprehensive numbers table

B1. Gate-A ON vs OFF ? the foundational decision (static, CONTROL vs `net\_crossings` ? 2')

variant	F1	P	R	seg-count ratio	# predicted segs
CONTROL (merge-all, no gate)	0.388	0.313	0.539	1.68	431
**Gate-A (`min_crossings` ? 2')	0.462	0.477	0.468	1.01	235

F1 = +0.074, 95% CI **[+0.042, +0.105]** (clears 0), sign test **6W / 0L, p = 0.031**.
Per-video ?F1 (all positive): adarsh +0.125, mahadevpura +0.117, testlarge +0.075, targetest +0.074,
gbaaddy +0.043, kushagra +0.009. **F1@k** all rise (F1@10 0.645?0.748, F1@25 0.603?0.698, F1@50 0.388?0.462). The over-seg ratio **1.68 ? 1.01** is the headline structural effect: Gate-A nearly eliminates the over-segmentation tIoU barely reflected, by pruning **196 of 431** low-crossing proposal windows (held-shuttle / single-feed / dead-time FPs ? the owner's bug). *Exactly reproduces the first-A/B archive (`docs/archives/quality-iterations/2026-06-first-ab-experiment/`).*

B2. Gate-B ON vs OFF (static, CONTROL vs `players\_in\_court` ? N')

Gate-B threshold	F1	?F1 vs CONTROL	95% CI	sign test	clears 0	# segs 431?
N ? 1	0.388	+0.000	[0, 0]	0W/0L	No	431 (no-op)
N ? 2	0.389	+0.000	[0, +0.001]	1W/0L	No	430 (?no-op)
N ? 3	0.307	**?0.081**	[?0.244, 0]	0W/1L	No	274
N ? 4	0.283	**?0.105**	[?0.305, +0.002]	1W/2L	No	231

Gate-B is inert-to-harmful on this corpus. At production-plausible thresholds (N?1/2) it is a
no-op ? nearly every window already shows ?1?2 detected persons because there are **no court polygons**, so `players\_in\_court` counts **spectators + adjacent-court players**, not just the two on the active court. Pushing N higher to force discrimination starts **deleting real rallies** (N?3: ?0.081; N?4: ?0.105). **Stacked on Gate-A, Gate-B adds exactly 0.000** (N?1 and N?2 both leave Gate-A's 0.462 unchanged). Gate-B is **not measurable as a win until court polygons exist**.

B3. Leave-one-signal-out ? marginal contribution to the learned 13-column fusion (LOVO)

(delta = F1(full 13 cols) ? F1(full minus that signal); positive ? the signal helps the scorer)

Full learned LOVO reference: **F1 = 0.366** (P 0.453 / R 0.341).

signal dropped	F1 full	F1 dropped	marginal ?F1	95% CI	sign test	clears 0
audio (3 cols)	0.366	0.286	**+0.079**	[?0.054, +0.203]	4W/2L (p=0.688)	No
`net_crossings` (as a soft column)	0.366	0.382	**?0.014**	[?0.083, +0.027]	4W/1L (p=0.375)	No
person (5 cols)	0.366	0.376	**?0.009**	[?0.074, +0.055]	2W/4L (p=0.688)	No
velocity block (peak+mean+active)	0.366	0.364	+0.002	[?0.010, +0.018]	2W/3L (p=1.000)	No
`duration`	0.366	0.367	?0.001	[?0.004, 0]	0W/2L (p=0.500)	No

Reading: the only signal carrying *positive* marginal information for the learned scorer is **audio** (+0.079; the scorer loses the most when audio is removed). **`net\_crossings` as a soft column is ~neutral-to-slightly-negative (?0.014)** ? its discriminative power is fully captured by the **hard Gate-A** form, not the learned column. Person is mildly negative (the no-court-polygon noise). Velocity/duration are ~zero marginal. **Every CI spans 0** (n=6) ? see the caveat.

B4. Incremental add-onto-shuttle (the P3/P4 framing, for cross-reference)

step	F1 A?B	?F1	95% CI	sign test	clears 0
shuttle ? +person (P3)	0.307	? 0.286	**?0.021**	[?0.165, +0.162]	2W/4L (p=0.688) No
shuttle ? +audio	0.307	? 0.375	**+0.068**	[?0.002, +0.140]	4W/2L (p=0.688) No
shuttle+person ? +audio (P4)	0.286	? 0.366	**+0.079**	[?0.054, +0.203]	4W/2L (p=0.688) No
shuttle ? +person+audio	0.307	? 0.366	+0.059	[?0.039, +0.156]	4W/2L (p=0.688) No

Consistent with the leaderboard (P3 ?0.021, P4 incr. +0.079). Audio is the lead learned cue; person is a wash on its own. Note **audio added directly to shuttle (+0.068, CI [?0.002, +0.140])** is the cleanest audio read (drops the person handicap) and its CI **barely** touches 0 ? the most promising near-promotable signal.

B5. Existing baselines (context ? IoU 0.5 vs human golden, from `eval\_baselines/` + QUALITY_ITERATIONS)

path	Mahadevpura (clean-ish)	Kushagra (multi-court)	corpus-mean F1
Heuristic (velocity windowing, LOVO n=6)	0.657	0.157	0.48
Gen-0 owned head (per-frame proto, LOVO n=6)	0.744	0.561	0.583
Two-stage WASB+Gemini refine	0.83	?	?
Gemini wrapper (`gemini-flash-latest`)	**0.93**	**0.66**	?

Wrapper baseline cost note: **well under \$1 for two videos** at flash pricing (serial, ~14 chunk requests). Multi-court drops the commodity wrapper 0.93 ? 0.66 ? that 0.66 is the owned model's ship target. **Caveat on cross-comparing rows:** the fusion/Gate numbers (\$B1?B4) are scored on a **recall-oriented proposal** candidate set (CONTROL ? 0.388 mean), a different operating regime than the heuristic/Gen-0/wrapper rows; don't read 0.462 vs 0.93 as "Gate-A loses to Gemini" ? they answer different questions (proposal-decision lift vs end-to-end detector quality).

C. Gate-A default flip ? implications (numbers AND foundations)

C1. What the numbers say

Gate-A is the **only** lever in this whole study that clears the honest promotion bar: ?F1 +0.074 with a CI that excludes 0 **and** a 6/0 sign test (p=0.031) ? it improves F1 on **every** golden match and slashes over-segmentation (1.68?1.01). It is cheap, GPU-free, deterministic, and was already corroborated on real footage (GX010125 offline re-score: 45 ? 27 rallies, 18 low-crossing FP windows dropped). By the repo's own "promote only on measured lift" rule, **Gate-A qualifies and Gate-B does not.**

C2. Extensibility ? does defaulting Gate-A ON constrain or enable future signal fusion?

It ENABLES, and here's the structural reason. Gate-A is a **proposal-stage precision filter**, not a verdict. It prunes structurally-impossible-to-be-a-rally windows (a held shuttle never crosses the net twice) **before** the learned scorer / Gemini ever sees them. That is strictly **complementary** to adding more signals:

- It **raises the floor** every later cue builds on. The full-13 learned fusion is measured on the **un-gated** CONTROL set (F1 0.366); a learned fusion measured **on top of a Gate-A-pruned set** starts from 0.462, with the worst FP class already removed ? every cue's job gets easier, not harder.

- It does **not** consume or mask any feature column. The persisted superset (shuttle+person+audio, 13 cols) stays intact in the JSONs; Gate-A only changes which **windows survive to scoring**. So audio (the promising +0.079 lead) and any future cue are added exactly as before ? Gate-A is orthogonal to the fusion layer, living one stage upstream.
- The harness already treats it as a first-class, swappable ``Variant`` knob (``min_crossings``), so a future cue that **subsumes** Gate-A's signal can be A/B'd against it and Gate-A flipped back off via config ? ****the experiment seam is preserved.****

****The one real constraint to name honestly:**** Gate-A is a ****hard gate****, not a soft feature. The leave-one-out numbers (\$B3) show the **soft** ``net_crossings`` column is ~neutral (?0.014) while the **hard** ``?2`` gate is the win ? i.e. the value is in the thresholding, which a learned scorer did ****not**** reproduce from the column at n=6. Defaulting the hard gate ON therefore bakes in a ****rule-based pre-filter alongside the learned path****. That is a deliberate two-stage shape (rule-gate ? learned/Gemini scorer), not pure end-to-end learning. It is reversible (config flag) and standard (proposal-filtering before scoring is textbook Temporal Action Localization), but it IS a foundational commitment to a ****hybrid rule+learned pipeline****, which the owner should make eyes-open.

C3. Scope of exploration ? what does the default flip open vs close?

- ****Opens:**** a clean, higher-precision substrate for the **next** signal hunt (audio first). Every future cue gets measured against the Gate-A floor, so a real lift is more likely to be a **real** lift and less likely to be "it cleaned up held-shuttle FPs the gate would have caught anyway."
- ****Opens:**** the held-shuttle / single-feed / dead-time FP class is closed structurally, freeing the exploration budget to chase the ****harder, temporal multi-court confound**** (the documented dead end for spatial/loudness heuristics ? only the learned temporal model cracks it).
- ****Closes nothing important:**** because the persisted superset is untouched, no future ablation is foreclosed. The only thing it "closes" is re-litigating held-shuttle FPs by other means.

C4. Foundational structure ? is it hard to reverse?

****No ? and this is the key de-risking fact.**** Gate-A as shipped is ``indexing.fusion.min_crossings`` (default 0 = OFF today). Flipping the default to 2 is a one-line config change; rollback is flipping it back (the repo's stated discipline: "rollback = flip a config flag, never ``git revert``"). It does not alter the DB schema, the persisted feature superset, or any registry contract. The **served-side** flip is the only part still owner-gated, and only because the production fusion segmenter (P2) is unverified on a live real-video run ? ****that's a verification gap, not a reversibility problem.****

The Gate-B side IS a foundational question, but the data says ****don't build on it yet****: it is inert without court polygons and harmful when forced. Committing to Gate-B now would bake a ****dependency on a court-mask source**** (manual polygon vs auto-homography ? an open owner decision) into the default path for ****zero measured benefit****. Keep Gate-B default-OFF; revisit only after court polygons exist, then

re-measure person *and* Gate-B together (the current ?0.021 person result is partly a
`court\_polygon=None`
artifact).

D. Recommendation framing (numbers + foundations, not numbers alone)

1. **Promote Gate-A to the default on the fusion/proposal path** (`indexing.fusion.min\_crossings = 2`).
- **Numbers:** the only lever that clears the honest bar (+0.074, CI clears 0, 6/0, p=0.031; over-seg 1.68?1.01) ? measured, reproduced, and corroborated on real footage.
- **Foundations:** it's an upstream, orthogonal, reversible precision pre-filter that **raises the floor** for every future cue without consuming any feature or foreclosing any ablation. The only eyes-open commitment is to a **hybrid rule+learned two-stage shape**, which is standard and config-reversible.
- **Gate on:** the live served-side flip stays owner-gated **only** until one real fusion-path run confirms P2 behaves as the offline re-score predicts (the GX010125 evidence already points the right way). This is verification, not a numbers question.
2. **Keep Gate-B default-OFF.** No measured benefit (no-op at N?2, harmful at N?3, +0.000 stacked on Gate-A). It is **blocked on court polygons**, not on more data. Do not bake a court-mask dependency into the default path for zero lift. Re-open Gate-B **and** the person cue together once polygons exist.
3. **Make audio the next promotion candidate, not person.** Audio is the strongest learned cue (leave-out +0.079; add-onto-shuttle +0.068 with a CI that only just touches 0). The honest blocker is n=6: every fusion CI spans 0. **Growing the golden corpus (Roora) is the highest-ROI unblock** ? it is the only thing that can turn the +0.07-ish audio lead and the person re-measurement into confirmable results.
4. **Be honest about the central caveat:** **at n=6 every learned-fusion CI spans 0.** Gate-A is the lone exception (its CI clears 0 and the sign test is significant) precisely because it's a strong, simple, structural rule that wins on all 6 videos ? not a marginal learned weight. So the numbers **decisively** support Gate-A and **decisively reject Gate-B**, but only **suggest** (cannot confirm) audio and **cannot exonerate** person until N grows. This is a data-influenced decision where the data is strong exactly where the foundational choice is being made (Gate-A) and weak where it isn't (the soft cues).

E. Is a served end-to-end Gemini run on the golden videos warranted? ? **NO (for the decision).**

Reason. The Gate-A/B decision lives in the **proposal/fusion** layer, upstream of Gemini, and is already settled by the persisted-feature ablation against golden labels (6/6, p=0.031, CI clears 0). A served Gemini run would measure a **different** thing ? final highlight-reel quality with the gate flipped on in the live path ? and it **cannot move a numbers-driven decision** because:

- There are **no golden labels** for new served-pipeline footage, so the output is divergence (``a_only`/`b_only`` windows), not F1 ? illustrative, not decisive.
- The golden videos' served behavior under Gate-A is **already corroborated** by the GX010125 offline re-score (45 ? 27 rallies, 18 low-crossing FP windows dropped, ~1.7 min), at zero Gemini cost.
- The remaining owner-gate on the served-side flip is a **verification** that the production fusion segmenter (P2) behaves as the offline path predicts ? a cheap single real-video run, **not** a golden-corpus Gemini sweep.

What a Gemini run WOULD be worth (optional, owner's call ? flagged, not run): a **confidence-builder** / demo*, not a decision input ? run the served fusion pipeline on ONE golden match with ``min_crossings`` OFF vs 2 and eyeball the two reels to confirm P2 drops the held-shuttle windows end-to-end.

- **Expected cost/time:** ~1 golden match ? tens of minutes wall-clock (GPU WASB/person pass + ffmpeg + serial Gemini chunks); Gemini spend **< \$1** at flash pricing (the wrapper baseline did two videos for under \$1). Two reels (gate OFF/ON) ? **2x that**, still sub-\$2.
- **Recommendation:** worth it once, as the **owner-gated served-side verification** of P2 (item D1's gate), but **not** as input to the Gate-A/B foundational decision ? the CPU ablation already decided that. Do not burn a multi-video Gemini sweep; it adds cost, not decision value.

Artifacts

- Runner: ``scratch/gate_ab_ablation.py`` · detail: ``scratch/gate_ab_detail.py``
- Raw measured result: ``scratch/gate_ab_ablation_result.json``
- Reuses: ``backend/eval/experiment.py``, ``backend/eval/eval_stats.py``, ``backend/eval/segmentation_metrics.py``, ``output/*_golden_features.json`` (the persisted substrate).
- Cross-refs: ``docs/QUALITY_ITERATIONS.md``, ``eval_baselines/experiment_leaderboard.json``, ``docs/archives/quality-iterations/2026-06-first-ab-experiment/README.md``.

41. user (2026-06-14T09:54:16.780Z)

File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\GATE_AB_NUMBERS_AND_IMPLICATIONS.md (file state is current in your context ? no need to Read it back)

42. assistant (2026-06-14T09:54:20.218Z)

Let me do a final consistency check ? verify the doc's headline numbers against the raw JSON, and confirm the `n_pred` pruning count (`431?235=196`) is right.

43. user (2026-06-14T09:54:24.330Z)

GateA dF1 0.074 CI 0.042 0.105 6 W 0 L p 0.031
 npred 431 -> 235 pruned 196
 full learned F1 0.366

44. assistant (2026-06-14T09:54:47.672Z)

All headline numbers confirmed (Gate-A ?F1 +0.074, CI [+0.042, +0.105], 6W/0L p=0.031, 196 windows pruned, full learned F1 0.366). The doc is consistent with the measured data. Everything was written only to ``scratch/``, tracked code was read-only, no ``--augment`` was needed (JSONs already had audio), and no Gemini/GPU was burned.

Report

I collected all golden-set numbers for the Gate-A/B foundational decision, ran a leave-one-signal-out ablation on the 6 persisted feature JSONs, and wrote the implications doc. All numbers are actually measured (CPU, seconds, reusing the shipped harness + honest stats). **Audio columns were already present in all 6 JSONs, so `--augment` was not needed.**

Deliverable: `**`C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\GATE_AB_NUMBERS_AND_IMPLICATIONS.md`**` (runner: ``scratch/gate_ab_ablation.py``; raw: ``scratch/gate_ab_ablation_result.json``).

Comprehensive numbers table (n=6, IoU?0.5, honest C1)

****Gate-A ON vs OFF (the foundational decision ? static):****

variant	F1	P	R	seg-ratio	#segs
CONTROL	0.388	0.313	0.539	1.68	431
Gate-A (?2)	**0.462**	0.477	0.468	**1.01**	235

?F1 ****+0.074****, CI ****[+0.042, +0.105]**** (clears 0), ****6W/0L, p=0.031****; 196 windows pruned; positive on all 6 videos. Exactly reproduces the first-A/B archive.

****Gate-B (players_in_court?N):**** no-op at N?2, ****harmful**** at N?3 (?0.081) / N?4 (?0.105); ****+0.000** stacked on Gate-A**. Inert without court polygons (counts spectators).

****Leave-one-signal-out marginal ?F1**** (full 13-col learned LOVO F1=0.366; delta = full ? dropped): audio ****+0.079****, net_crossings-as-column ****?0.014****, person ****?0.009****, velocity +0.002, duration ?0.001 ? ****every CI spans 0****.

****Incremental:**** +person (P3) ?0.021; +audio +0.068 (CI [?0.002,+0.140]); +audio on person (P4) +0.079.

****Existing baselines:**** heuristic 0.48 mean / Gemini wrapper 0.93 clean?0.66 multi-court / Gen-0 0.583.

Key implications

- ****Gate-A is the only lever clearing the honest bar.**** It's an ****upstream, orthogonal, reversible**** precision pre-filter (one-line ``indexing.fusion.min_crossings=2``, default-OFF today; rollback = flip the flag). It does ****not**** consume any feature column or foreclose any ablation ? the persisted 13-col superset stays intact ? so it ****raises the floor**** every future cue builds on. The one eyes-open commitment: the ***hard gate*** (not the soft ``net_crossings`` column, which is ~neutral at ?0.014) bakes in a ****hybrid rule+learned two-stage shape**** ? standard and config-reversible.
- ****Gate-B is a foundational trap right now:**** committing to it bakes a court-mask-source dependency into the default for ****zero measured benefit****. Keep default-OFF; re-open with person once court polygons exist.
- ****Central caveat (honest):**** at n=6 every ***learned-fusion*** CI spans 0. Gate-A is the lone exception precisely because it's a strong structural rule winning 6/6, not a marginal learned weight. Audio is the least soft cue but only ***suggestive***; growing the corpus (Roora) is the highest-ROI unblock.

Recommendation framing

Promote ****Gate-A**** (numbers decisive + foundations enabling/reversible); keep ****Gate-B OFF**** (no lift, blocked on polygons); make ****audio**** the next promotion candidate over person; be explicit that the data is strong exactly where the foundational choice is made (Gate-A) and weak where it isn't (soft cues).

Served Gemini run ? ****NO**** (for the decision)

The Gate-A/B decision is in the proposal/fusion layer ****upstream of Gemini**** and already settled against golden labels (6/6, p=0.031). A served run can't move a numbers decision (no golden labels on new footage ? divergence, not F1) and the served behavior is already corroborated by the GX010125 offline re-score (45?27 rallies, 18 FP windows dropped, zero Gemini cost). Flagged for the owner as an ****optional one-match confidence-builder / P2 served-side verification**** only: ~tens of minutes wall-clock, Gemini spend ****<\$1 per reel**** (sub-\$2 for an OFF/ON pair) ? worth it once as the owner-gated served verification, not as input to this decision. Do not burn a multi-video Gemini sweep.