

claude-windows / 6a82308f-69d6-43a8-9ea4-d8eaf497253d

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6a82308f-69d6-43a8-9ea4-d8eaf497253d\subagents\agent-a41d134f20a2a7d97.jsonl`  
- SHA-256: `d66a8731221a5c377f8c87b178a45939e90d11592e50a02c4b473cd8dcd9cc06`  
- Source modified: `2026-06-16T03:13:38+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `subagents`  
- session_id: `6a82308f-69d6-43a8-9ea4-d8eaf497253d`
```

Transcript

1. user (2026-06-16T03:12:12.068Z)

Survey the rally-detection QUALITY / HARNESS / strategy design docs in this repo (cwd = the sports highlight indexer) so a new foundational research doc can build ON them and cross-link, not duplicate.

Read and summarize each (skip gracefully if absent): docs/README.md, docs/CODE_MAP.md, docs/EXPERIMENT_HARNESS.md, docs/MULTI_SIGNAL_FUSION_PLAN.md, docs/ANALYZERS_AND_RECONCILERS.md, docs/HOW_RALLY_DETECTION_WORKS.md, docs/QUALITY_ITERATIONS.md, docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md, docs/NEXT_STEPS.md, docs/ROORA_LABELING_GUIDELINES.md, docs/PERSONALIZATION_PLAN.md, and anything under docs/archives/quality-iterations/.

For EACH doc report: (1) what it covers, (2) its STATUS (designed-only / shipped / owner-gated / not-started ? look for explicit status markers), (3) the key decisions/claims/numbers.

Then answer precisely:

- **CRUCIAL:** What did ANALYZERS_AND_RECONCILERS.md §13 (the multi-agent literature review) already cover? List the specific patterns surveyed and the verdicts (e.g. "candidate-windows+scorer = TAL", "reconciler = TAS post-proc", "signal-not-verdict = late/score-level fusion", which were rated novel vs standard prior art). A new research effort must go BEYOND this ? so capture exactly what's already surveyed + cited.
- What are the prior MEASURED results recorded (e.g. Gate-A ?F1 +0.074 6/6 p=0.031; audio +0.079; person ?0.021; the learned-variant/threshold-tuning numbers; shuttle-only F1?0.307; LOVO 0.626)? Capture the numbers + their context.
- What is the stated eval approach (LOVO/LOGO, IoU threshold, honest stats), the eval?serve lesson, and the "signals?learned-scorer, NO special-casing" spine directive?
- What does the owned-model plan say about the base model + licensing constraints (Qwen3-VL/InternVL3; MIT/Apache/BSD only; paid LLMs inference-only never training data; exclude minors)?

Return structured markdown, distilled, with doc + section anchors. NO full dumps.

2. assistant (2026-06-16T03:12:16.569Z)

I'll systematically survey the documentation in the sports highlight indexer repo to extract the key design decisions, measured results, and strategic direction.

3. user (2026-06-16T03:12:16.717Z)

```
docs\GOLDEN_SET_IMPLEMENTATION.md  
docs\GOLDEN_SET_PHASE1_VIDEOS.md  
docs\PMF_MARKET_RESEARCH.md  
docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md  
docs\HOSTING_PLAN.md  
docs\UI_PROPOSAL.md  
docs\UI_ALPHA_EXECUTION_PLAN.md  
docs\FIELD_VISIT_ANSWER_SHEET.md
```

docs\FIELD_VISIT_PLAYBOOK.md
docs\I18N_PLAN.md
docs\PMF_FIELD_SIGNALS.md
docs\PRODUCT_FIELD_ASSOCIATE_FLYER.md
docs\VIDEO_LOCALITY_MODEL.md
docs\GOLDEN_SET_VENDORING.md
docs\HOW_RALLY_DETECTION_WORKS.md
docs\STORAGE_SHARING_MODEL.md
docs\PER_VIDEO_OUTPUT_LAYOUT.md
docs\QUALITY_ITERATIONS.md
docs\ABLATION_LAYER.md
docs\ANALYZERS_AND_RECONCILERS.md
docs\BACKLOG.md
docs\CODE_AUDIT_AND_TEST_HARDENING.md
docs\COMMERCIALIZATION.md
docs\DEFERRED_HARDENING_PLAN.md
docs\EVALUATION.md
docs\EXECUTION_POLICY.md
docs\EXPERIMENT_HARNESS.md
docs\GOLDEN_DATA_SHARING.md
docs\HARDENING_LOOP_HANDOFF.md
docs\MULTI_SIGNAL_FUSION_PLAN.md
docs\NEXT_STEPS.md
docs\PERSONALIZATION_PLAN.md
docs\PLATFORM_ARCHITECTURE.md
docs\POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md
docs\README.md
docs\ROADMAP.md
docs\ROORA_LABELING_GUIDELINES.md
docs\TRACKNET_WSL_SETUP.md
docs\VIDEO_RECORDING_GUIDELINES.md
docs\KHELSUTRA_PRODUCT_OVERVIEW.md
docs\CODE_MAP.md

4. user (2026-06-16T03:12:16.805Z)

docs\archives\COMMERCIAL_READINESS_REVIEW.md
docs\archives\MONETIZATION_AUDIT.md
docs\archives\checkpoints\2026-06-eval-baseline\README.md
docs\archives\checkpoints\2026-06-logging-reporting\README.md
docs\archives\checkpoints\2026-06-scaffolding-complete\README.md
docs\archives\past_projects\README.md
docs\archives\research\AUDIO_E1_FINDINGS.md
docs\archives\research\AUDIO_RALLY_DETECTION_PLAN.md
docs\archives\research\RALLY_DETECTION_RESEARCH.md
docs\archives\checkpoints\2026-06-strategy-foundation\README.md
docs\archives\quality-iterations\2026-06-human-gt-tuning\README.md
docs\archives\quality-iterations\2026-06-initial\README.md
docs\archives\quality-iterations\2026-06-multivideo-lovo\README.md
docs\archives\research\OWNED_MODEL_TRAINING_STUDY.md
docs\archives\past_projects\low-regret-moves\LOW_REGRET_MOVES_PLAN.md
docs\archives\decisions\DECISIONS.md
docs\archives\TEST_RUN_SUMMARY.md
docs\archives\quality-iterations\2026-06-first-ab-experiment\README.md
docs\archives\quality-iterations\README.md
docs\archives\checkpoints\2026-06-14-hardening-loop-complete.md
docs\archives\CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md
docs\archives\DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md
docs\archives\HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md
docs\archives\README.md

5. assistant (2026-06-16T03:12:19.744Z)

Now I'll read the key docs you specified:

6. user (2026-06-16T03:12:20.089Z)

```
1      - [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Latest (v8) vendor
labeling guidelines mirrored from the collector's `docs/annotation/`, with the fusion-driven
court-calibration / singles-doubles additions marked `? ADDED (fusion)`.
2      - [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
3      - [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (P1 SHIPPED) for A/B + shadow
+ SxS experimentation on rally-detection variants with zero production regression ?
experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift,
rollback = flip a flag). Composes the existing `backend/eval/` machinery (`calibration`,
`regression`, `metrics`); implemented in `backend/eval/experiment.py`.
4      - [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) ? Living log of rally-detection
quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-
chronological record; terminal-state snapshots get archived under `archives/quality-
iterations/`.
5      - [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated,
NOT started): a signal-fusion framework for rally detection ? producers emit confidence-
scored signals at three temporal scopes (frame detectors / window analyzers / session
analyzers), a learned scorer weights them, and an auditable report-first reconciliation
pass lints the decision against the evidence. The spine: no special-casing in the decision
path (signals?scorer, never `if/else`). Worked examples: a service-fault analyzer, a warm-
up/practice span detector, and consuming a per-frame shuttle-state detector. A peer to
`EXPERIMENT_HARNESS.md` and `MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).
6      - [PERSONALIZATION_PLAN.md](PERSONALIZATION_PLAN.md) ? Design (owner-adopted hybrid;
Phase-0 landing) for the per-user personalization feature on a generalization-first
baseline: a thin per-user tuning layer gated by the honest small-N statistics (a min-N
promotion floor + a structural-guard exemption, never a service gate) + a contribution
flywheel (free-tier videos pool into the owner-trained baseline; "a tool that gets better the
more you help it"). Records the locked owner decisions (identity, storage, min_n
semantics, Gate-A default, free-tier pooling). Built so far: per-user promotion gate (#141) +
held-out-USER learning curve (#142).
7
8      ## Practical / ongoing
9      - [TRACKNET_WSL_SETUP.md](TRACKNET_WSL_SETUP.md)
10     - Golden set docs (`GOLDEN_SET`)
11     - Evaluation harness and calibration drivers (in `backend/eval/`)
12     - [GOLDEN_DATA_SHARING.md](GOLDEN_DATA_SHARING.md) ? How the GPU box and CPU dev/agent
sessions share golden artifacts (trajectory CSVs, `_golden_features.json`) via a OneDrive
folder; the GPU-extract ? CPU-litmus split. See #175.
13
14     ## Historical documents
15     Completed plans, old research, ADRs, and checkpoints live in [`archives/`](archives/).
16
17     - [CODE_AUDIT_AND_TEST_HARDENING.md](CODE_AUDIT_AND_TEST_HARDENING.md) ? COMPLETED
(2026-06-14; T1?T5 + full loop + verification per #162). Living record of the June 2026 plumbing
/ contract / architectural audit, 626-test infrastructure, positive test plans for the 5
identified issues, multi-layer regression protocol (including ablation/experiment layer), and
pre/post-change verification checklist. See `docs/archives/CODE_AUDIT_AND_TEST_HARDENING-
COMPLETED-2026-06-14.md` and `docs/archives/checkpoints/2026-06-14-hardening-loop-complete.md`
for the final archived state. (Top-level is now pointer stub per review; full content in
archives/.)
18
19     See `archives/README.md` for the umbrella policy and `archives/past_projects/README.md`
for classification of small tech debt / starter projects.
```

7. user (2026-06-16T03:12:20.515Z)

```
1      # Experiment harness ? A/B + shadow eval for rally-detection quality (no-regression by
design)
2
3      Status: P1 IMPLEMENTED (`backend/eval/experiment.py` +
`tests/test_experiment.py`, shipped
4      2026-06-13) ? the `Variant` layer, `compare()` (labeled A/B, LOVO-honest),
`compare_unlabeled()`
5      (SxS on NEW unlabeled footage), and the leaderboard. P2?P4 remain owner-gated.
```

Phases below are

```

6      sequenced but each is a separate PR.
7
8      > **Companion docs:** the cues this experiments *on* are in
9      > [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md); the scoring mechanics are
10     > [EVALUATION.md](EVALUATION.md); the no-regression gate + golden labels +
`baseline.json` come from
11     > [GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md); the modular-cue architecture is
**ADR-007**
12     > ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the distilled
classifier is
13     > **ADR-004**. The 2-layer mental model is
[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md).
14
15     ---
16
17     ## 1. Context ? why this, why now
18
19     PR #112 (merged) added the multi-signal fusion design ? shuttle + person + audio cues,
all **default
20     OFF**. The owner's directive for *how* we evolve quality from here:
21
22     > **"Evolve the codebase so we can keep finding new ways to improve quality and
experiment with them ?
23     > or add them as a signal ? without suffering any regressions if the idea doesn't work
out. Rolling
24     > back the code would be too dangerous."**
25
26     So **experimentation and deployment must be decoupled.** The contract:
27
28     1. Merge experimental code freely ? it's **gated OFF**, so it cannot change production
behavior.
29     2. Test it **offline against golden labels** (most experiments never touch the served
path at all).
30     3. Promote a variant to default **only on measured lift**, through a CI eval gate.
31     4. **"Rollback" = flip a config flag, never `git revert`. ** The proven path is a pinned
variant that is
32     always present.
33
34     **The pleasant surprise (Explore audit, 2026-06-13):** ~80% of this harness already
exists and is
35     simply not composed *as a system*. The missing piece is a thin **variant/experiment
layer + the
36     discipline**, not a platform. §4 is the exact code surface so the next session does not
re-discover it.
37
38     ---
39
40     ## 2. The thesis ? separate expensive DETECTION from cheap DECISION
41
42     The expensive, GPU/WSL-bound, risky part is **detection**: "what signals exist per
frame" (shuttle
43     WASB trajectory, person tracks/boxes, audio hits). The cheap, swappable, pure-Python
part is the
44     **decision**: "given those signals, where are the rallies" (windowing + gates + the
classifier).
45
46     > **Run detection once per video, persist all per-cue signals + per-candidate features,
then re-score
47     > any number of variants OFFLINE ? deterministically, with no GPU and zero production
risk.**
48
49     This is not aspirational ? `distill_local.py` and `calibrate_local.py` already do
exactly this on
50     stored WASB trajectories. The candidate-segments table (`stats` JSON) is the persistence

```

```

substrate.
51     ? Most experiments are pure re-scoring of stored data; they never enter the served
pipeline.
52
53     ---
54
55     ## 3. What already exists (the audit ? do NOT rebuild)
56
57     Exact public surface, with `file:line`. **This table is the anti-duplicate-work
artifact** ? start here.
58
59     ### 3.1 Scoring (variant-agnostic ? grades any `List[Interval]`)
60     `backend/eval/metrics.py`
61     - `interval_iou(a, b) -> float` (:15) ? temporal IoU of two `(start,end)` intervals.
62     - `match_intervals(preds, gts, iou_threshold=0.5) -> MatchResult` (:48) ? greedy 1-to-1
IoU matching.
63     - `score(preds, gts, iou_threshold=0.5, match_result=None) -> Score` (:104) ? P/R/F1,
mean IoU, boundary errors; `Score.as_dict()``.
64     - `pr_curve(preds, gts, thresholds=None) -> List[Score]` (:138).
65
66     ### 3.2 A/B + cross-validation primitives (the comparison engine already exists)
67     `backend/eval/calibration.py`
68     - `improvement_report(predict_fn, baseline_config, tuned_config, gts, iou_threshold=0.5)
-> dict` (:148) ? **before/after metrics + %?. This IS the A/B delta.**
69     - `leave_one_video_out(videos, configs, metric="f1", iou_threshold=0.5) -> dict` (:113)
? **LOVO: pick on other videos, score on held-out video (honest cross-video).**
70     - `cross_validate(predict_fn, configs, gts, k=5, metric="recall", iou_threshold=0.5) ->
dict` (:72) ? within-video k-fold overfit guard.
71     - `select_best(predict_fn, configs, gts, metric="f1", iou_threshold=0.5) -> (Config,
float)` (:61).
72     - `evaluate(preds, gts, iou_threshold=0.5) -> Score` (:41); `kfold_indices(n, k)` (:46).
73     - LOVO video dict shape: `{"name": str, "predict_fn": PredictFn, "gts": [Interval]}`.
74
75     ### 3.3 No-regression gate + baselines (the promotion gate already exists)
76     `backend/eval/regression.py`
77     - `regress(db, video_id, ground_truth, stage="final", iou_threshold=0.5, detector=None,
tolerance=DEFAULT_TOLERANCE, baseline_path=DEFAULT_BASELINE_PATH) -> RegressionResult` (:104) ?
score stored preds vs GT, flag if P **or** R drops > tolerance.
78     - `regress_with_provider(...)` (:160) ? same, fetching GT via the annotation provider.
79     - `save_baseline(video_id, stage, precision, recall, f1, ..., gt_fingerprint=None)`
(:80) . `load_baselines(path) -> dict` (:68).
80     - `RegressionResult` (:43): `regressed`, `reasons`, `gt_hash`, baseline P/R,
`tolerance`, `has_baseline`; `as_dict()`.
81     - Default baseline file: `eval_baselines/regression_baseline.json` (:33).
82
83     `run_regression.py` ? CI-gatable CLI, `main(argv)` (:52). **Exit codes: 0 ok / 1
regression / 2 no-DB / 3 no-baseline.** Flags: `--video-id --tier --annotations
--stage{candidates,final} --detector --iou --tolerance --baseline --update-baseline --allow-
missing-baseline --json`.
84
85     ### 3.4 Pinnable model artifact
86     `backend/eval/classifier.py` ? `LogisticRegression` (pure numpy, L2):
`fit(X,y,feature_names=,class_weight="balanced")` (:40), `predict_proba` (:75),
`predict(threshold=0.5)` (:81), `to_dict`/`from_dict` (:86/:97), `save`/`load` (JSON)
(:107/:111). **A trained model = one JSON file ? a pinnable, hashable variant artifact.**
87
88     ### 3.5 Offline re-scoring substrate (persist once, re-score forever)
89     `backend/infrastructure/database.py`
90     - `add_candidate_segment(video_id, detector, start_time, end_time, chunk_index=None,
confidence=None, stats=None, detection_params=None) -> Optional[int]` (:318) ?
`stats`/`detection_params` stored as JSON; `candidate_segments` table cols incl. `stats`,
`detection_params`, `human_label`, `confirmed_segment_id` (schema :97).
91     - `query_candidate_segments(video_id, detector=None, only_unlabeled=False) ->
List[dict]` (:355) ? returns rows with deserialized `stats`.
92

```

```

93     ### 3.6 Label assignment + feature glue
94     `backend/eval/training.py`
95     - `label_candidates(candidate_intervals, gt_intervals, iou_threshold=0.5) -> List[int]`
96     (:50) ? y-labels by IoU match to golden (same matcher as the evaluator).
97     - `build_training_set(candidate_rows, gt_intervals, iou_threshold=0.5,
98     feature_fn=extract_yolo_features) -> (X, y, intervals)` (:64) ? pluggable `feature_fn`.
99     ### 3.7 Existing LOVO drivers to mirror (not duplicate)
100    - `backend/eval/distill_local.py` ? `run(specs, iou=0.5, threshold=0.5, model_out=None)`
101    (:98), CLI `python -m backend.eval.distill_local --manifest videos.json --model-out ...`;
102    `FEATURE_NAMES = [duration, peak_velocity, mean_velocity, active_ratio, net_crossings]` (:40);
103    already does LOVO over Gemini silver labels.
104    - `backend/eval/calibrate_local.py` ? `run(specs, iou=0.5, metric="f1")` (:77);
105    `LocalConfig = (velocity_thresh, merge_gap, min_window_duration, min_crossings)`; grid + LOVO.
106    - `backend/eval/calibrate_wasb.py` ? `run(...)` , `load_golden_gts()`;
107    `backend/eval/gt_loader.py::load_gt_intervals(csv_path, *, strict=True)` is THE canonical golden
108    reader.
109    ### 3.8 Registries + config knobs (variant selection)
110    - `backend/pipeline/segmenters/base.py` ? `SegmenterRegistry` (:35), singleton
111    `segmenter_registry` (:63); registered (6): `motion`, `gemini`, `yolo_hybrid`,
112    `tracknet_hybrid`, `wasb_hybrid`, `fusion_hybrid`.
113    - `backend/annotations/base.py` ? `AnnotationProviderRegistry`, singleton
114    `annotation_registry`; providers `file`, `auto`.
115    - `backend/config/models.py::IndexingConfig` ? `default_segmenter` (:65, default
116    `"yolo_hybrid"`), `detector_model` (:72), `detector_score_threshold` (:73),
117    `yolo_velocity_threshold` (:74), `yolo_frame_skip` (:75), `min_segment_duration` (:68),
118    `dead_time_merge_gap` (:69), `motion_threshold` (:67). The `min_crossings` knob now lives on
119    `FusionConfig` (default `0` = OFF; also `min_players`, both served Gate A/B thresholds ? see
120    fusion P2).
121    ### 3.9 Confirmed ABSENT (no duplication risk)
122    Explore found no existing experiment / variant / shadow / A-B machinery ? only a
123    stray "A/B test"
124    aspiration comment in `backend/pipeline/detectors/base.py` and an "experiment E1" note
125    in
126    `archives/research/AUDIO_RALLY_DETECTION_PLAN.md`. The `regress()` + baseline path is
127    the canonical comparison mechanism.
128    ---
129    ## 4. The variant abstraction (the one thin thing to add)
130    A Variant = a declarative recipe, not a code branch:
131    ---
132    Variant = (segmenter_name, IndexingConfig overrides, enabled cues + per-cue params,
133    classifier model path/hash)
134    ---
135    JSON-serializable; selected via the existing `segmenter_registry` +
136    `IndexingConfig.default_segmenter`.
137    Control = a pinned, frozen variant (recipe + model hash) that is always registered
138    and is the
139    default. Adding a new cue = a new Variant differing by one flag ? the control recipe is
140    untouched.
141    ---
142    ## 5. Three experiment modes
143    ### 5.1 Offline A/B (primary, label-driven, zero production risk)
144    `compare(videos, variant_A, variant_B)` ? a thin driver composing existing
145    functions:
146    `query_candidate_segments()` (load persisted features) ? re-score each variant ?

```

```

`calibration.improvement_report()`
134 + `calibration.leave_one_video_out()` + `metrics.score()`, graded vs golden
(`gt_loader.load_gt_intervals`).
135 Reports per-video **and** LOVO-honest aggregate IoU/P/R/F1 deltas. **This is the A/B
test, and it needs
136 almost no new scoring code** ? only the glue.
137
138 **Honest reporting (default ON, additive):** `compare(..., honest_stats=True)` also
attaches a `"honest"`
139 block ? per-metric bootstrap **CIS**, a paired **?F1 CI + exact sign test** (C1), and
**F1@k +
140 segment-count ratio** (C4) ? *alongside* the bare-mean fields (existing output
unchanged; `honest_stats=False`
141 restores the legacy shape). A bare LOVO mean at n?6 is not promotable on its own; the
lift is real when the
142 ?F1 CI clears 0 **and** the sign test agrees. Wiring of the course corrections in
143 [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) §13/C1+C4 (helpers:
`backend/eval/eval_stats.py`,
144 `segmentation_metrics.py`); `record_experiment` carries the honest ?F1 into the
leaderboard.
145
146 ### 5.2 Shadow (for cues that touch detection, e.g. the person detector)
147 The new cue runs and **persists its signal** into the candidate
`stats`/`detection_params`, but the
148 *served* decision stays control. Offline you then ask "what would variant B have
scored?" via §5.1 ?
149 B never affects output until it wins. This is how a detection-touching change (not just
a re-scoring
150 change) earns its way in without risk.
151
152 ### 5.3 Promotion gate (no-regression, CI-enforced)
153 `regression.regress()` + `baseline.json` + `run_regression.py` exit codes. A variant
becomes the new
154 default **only if** it doesn't regress the golden set beyond `tolerance` (exit 1
blocks). Control is
155 pinned ? **rollback = flip `default_segmenter`/variant config, never a code revert.**
156
157 ---
158
159 ## 6. No-regression guarantees (the heart of the owner's ask)
160
161 - **New cues/signals default OFF** ? per-cue enable flag in `IndexingConfig`; merged
code can't change
162 behavior until explicitly enabled.
163 - **Control variant pinned** ? frozen recipe + classifier model hash, always registered
& default.
164 - **Promotion only via the eval gate** ? `run_regression.py` in CI (exit 1 blocks); no
silent default change.
165 - **Experiments are records** ? variant-spec hash ? per-video + LOVO scores + label-set
hash + timestamp,
166 persisted to a small JSON **experiment leaderboard** (generalizes `baseline.json`).
Answers "which
167 signals ever helped, on which video slices," and makes every result reproducible.
168 - **Decoupled events** ? "merge an experiment" and "change the default" are separate,
independently
169 revertible actions.
170
171 ---
172
173 ## 7. What to build later (gap list ? thin, additive, owner-gated)
174
175 1. ~~~`backend/eval/experiment.py`~~ ? a `Variant` dataclass + `compare(videos, a, b)`
composing the §3
176 functions (no new scoring math).~~ **DONE (2026-06-13).** Shipped `Variant` (+ pinned
`CONTROL`),

```

```

177     `compare` (labeled A/B with LOVO-honest learned-variant training, mirroring
`distill_local`),
178     `compare_unlabeled` (the SxS-on-new-videos mode ? agreed / A-only / B-only via
`match_intervals`,
179     no GT), `record_experiment` (? `eval_baselines/experiment_leaderboard.json`), and
`load_video_candidates`
180     (the `query_candidate_segments` bridge). JSON variant recipes first as recommended; a
`VariantRegistry`
181     only if the count grows. 19 tests incl. the no-regression invariant.
182     2. Persist per-cue features into candidate `stats` (person: `players_in_court`,
`two_sided_motion`,
183     `in_court_ratio`, `shuttle_player_colocation`; shuttle-state: `net_crossings`,
`direction_reversals`;
184     audio: hit times) so any future variant re-scores offline ? ties to
`MULTI_SIGNAL_FUSION_PLAN.md` §3.2/§4.
185     3. Per-cue enable flags in `IndexingConfig` (default `False`); `default_segmenter`
already exists.
186     Add the `min_crossings` knob (fusion P2) here too.
187     4. CI promotion lane using `run_regression.py` against the golden `baseline.json`;
write the
188     experiment-leaderboard JSON.
189
190     ---
191
192     ## 8. Phasing (owner-gated; ONE PR per slice, off `master`)
193
194     - P0 ? this doc, indexed in `docs/README.md`, cross-linked from
`MULTI_SIGNAL_FUSION_PLAN.md` and
195     `GOLDEN_SET_VENDURING.md`. *(this deliverable; docs-only, no code.)*
196     - P1 ? `experiment.py` `compare()` ? DONE (2026-06-13): thin wrapper over
`calibration` +
197     `metrics` + `query_candidate_segments`; pure-logic, fully tested in this dev container
(`compare`,
198     `compare_unlabeled`, `record_experiment`, the no-regression invariant).
`compare_unlabeled` adds the
199     SxS read on new unlabeled footage the owner asked for (run control + candidate on
a fresh video,
200     diff where they agree/diverge ? two reels via the existing `VideoCompiler` is the GPU-
side follow-on).
201     - P2 ? per-cue enable flags + feature persistence into candidate `stats` (default
OFF) ? couples
202     with fusion P1/P2.
203     - P3 ? CI promotion gate: `run_regression.py` lane on the golden `baseline.json`;
experiment-leaderboard JSON.
204
205     ---
206
207     ## 9. Scope guardrail
208
209     Single-tenant / local-first (per `PLATFORM_ARCHITECTURE.md`): a JSON experiment
registry + the
210     existing harness ? explicitly NOT MLflow / W&B / a service. Keep it the size of
`baseline.json`.
211
212     ---
213
214     ## 10. Verification
215
216     - Offline, runnable in this dev container (no GPU/torch needed for the glue): unit
test that
217     `compare()` is deterministic on synthetic candidate rows; a no-regression invariant
test asserting
218     "all new-cue flags OFF ? predictions byte-identical to control"; a CI-gate test
asserting
219     `run_regression.py` exit codes (0 ok / 1 regress / 3 no-baseline).

```



```

220 - **Machine-side (owner, GPU/WSL):** persist candidates once on a real clip, run
`compare(control,
221 fusion)` LOVO, confirm the reported lift matches a manual `metrics.score()` spot-
check.
222 - **Docs (P0):** links resolve; indexed in `docs/README.md`; cross-linked from
223 `MULTI_SIGNAL_FUSION_PLAN.md` and `GOLDEN_SET_VENDORING.md`.
224 </content>
225

```

8. user (2026-06-16T03:12:20.724Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-151 of 379 total (37029 tokens, cap 25000). Call Read with offset=152 limit=151 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-reminder>

```

1      # CODE_MAP.md
2
3      **Purpose:** cold-start navigation map for the badminton-highlight-indexer ? a self-
hosted video ? highlight-reel pipeline (CV/AI rally detection ? SQLite ? ffmpeg stitch) with a
segmenter-agnostic eval/calibration stack, typed config, and per-video observability reporting.
4
5      **Regenerate:** rerun the `build-code-map` workflow (parallel subsystem readers ? this
doc). Re-run after any registry/contract/WSL-interface change, or after a directory restructure
(this revision follows #40).
6
7      ? **2026-06-10 audit-hardening sweep changed several behaviors documented below ? trust
git history over stale lines until the next full regen:** `add_video` is now an **UPSERT** (no
FK-cascade wipe on re-ingest; segmentation is destroy-AFTER-confirm via
`segment_watermarks`/`prune_segments`); detector artifacts are keyed by **video_id**, not
basename; **`indexing.wasb.*` is now a real typed config** (was silently dropped); `timeout_sec`
defaults **1800s** (was 0/hang-forever); SQLite opens with **WAL + busy_timeout**;
```

`match\_intervals` requires real overlap (iou>0); `run\_regression` baseline lives in tracked
`eval\_baselines/` and **exits 3 on a missing baseline**;

the API validates
`video\_path`/`output\_filename`. New tests: `test\_api\_endpoints`, `test\_compiler`,
`test\_wasb\_runner`, `test\_reingest\_safety`, `test\_detector\_keying`, `test\_config\_wasb`,
`test\_path\_safety`, `test\_reliability\_hardening` (+ collector `test\_licensing\_gate`,
`test\_import\_guard`, `test\_cvat\_robustness`).

```

8
9      **LEGEND**
10     - `@path` = file (always repo-relative)
11     - `::sym` = symbol (class/func/const) in the nearest preceding `@path`
12     - `->` = dataflow / call / "produces"
13     - `$` = data contract (canonical shape)
14     - `?` = gotcha / footgun
15     - `?` = registry / plugin extension point
16     - `WSL` = runs inside WSL2 conda env, NOT the Windows Python process
17
18     ---
19
20     ## 0. REPO LAYOUT ? where new files go (READ BEFORE committing a new file)
21
22     Top-level tree and the **placement rule** for each kind of file. When unsure, prefer an
existing
23     **registry / extension point** (`?`) over a new top-level file.
24
25     | You're adding? | Put it in? | Notes |
26     |---|---|---|
27     | Backend / runtime code | `backend/<subsystem>/` |
`pipeline/{segmenters,detectors,validators}`, `providers/`, `sports/`, `annotations/`,
`infrastructure/`, `reporting/`, `stitcher/`, `utils/`, `config/`, `features/`, `api/`. Most
additions are **register a class in the relevant `?` registry**, not a new module tree. |
28     | **Eval LIBRARY** (importable, pure ? metrics, calibration, a scorer, a feature) |
`backend/eval/` | Pure core, no side-effects at import; CLI/I-O lives in a separate driver.
These are **imported widely** (e.g. `calibrate_wasb` by 8 sites) ? they are library, not

```

```

scripts; do NOT move them out. |
29 | **Standalone run-driver / one-off entry CLI** (has `__main__`, NOT imported by app
code) | **`scripts/`** | e.g. `scripts/run_phase3_eval.py`, `scripts/run_process_and_stitch.py`.
Run as `python scripts/<name>.py`. If it needs `load_dotenv()`/`sys.path` before importing
backend, add `scripts/<name>.py" = ["E402"]` to `ruff.toml`. |
30 | **Ops / maintenance tool** (not part of the served app ? cache GC, batch admin) |
`tools/` | Run as `python -m tools.<name>` (it's an importable package). e.g.
`tools/cleanup_caches.py`. Keep the destructive decision path pure + unit-tested. |
31 | Training code (model fitting) | `training/` | Behind the CI-enforced import wall
(ADR-008) ? `backend.main` must not import it. Own `requirements-train.txt`. |
32 | A test | `tests/test_*.py` | pytest auto-discovers; the full-suite lane has a **70%
coverage floor**. Pure/offline + mocked (no GPU/WSL/network). |
33 | A doc (design, plan, living log) | `docs/` | Index it in `docs/README.md`. Only
**terminal-state** (DONE/REJECTED) work moves to `docs/archives/`. |
34 | A throwaway repro / debug script | `scratch/` | **Gitignored** + excluded from
ruff/pytest. Never imported by app code. |
35 | A model artifact (weights + manifest) | `artifacts/<name>/<ver>/` | ADR-008:
**manifest committed** (provenance), **weights gitignored**. |
36 | Eval baseline / leaderboard JSON | `eval_baselines/` | Tracked (survives a clean
checkout); the no-regression gate reads it. |
37 | Pipeline run outputs (reels, trajectories, DBs, frames) | `output/` | **Gitignored.**
Per-video isolation is planned in `PER_VIDEO_OUTPUT_LAYOUT.md`. |
38
39 **Canonical files that stay at the repo ROOT** (referenced by CI / packaging / docs ? do
not move):
40 `pyproject.toml` (PEP 621 packaging, hatchling; `indexer` entry point) .
`run_all_tests.py` (CI: `ci.yml`) .
41 `run_regression.py` (the no-regression gate CLI, cited in `EXPERIMENT_HARNESS.md`) .
`run_eval.py`
42 (grandfathered ? imported by `tests/test_eval_harness.py`). The old diagnostic
`setup.py` moved to
43 `scripts/doctor.py` (it shadowed the build system); invoke via `python
scripts/doctor.py` or `indexer --setup`.
44 New standalone run-drivers go in **`scripts/`**, not the root.
45
46 ---
47
48 ## 1. PIPELINE
49
50 ### 1A. Runtime: video file -> highlight reel
51 ```
52 typed config load (built-in defaults -> config.json overrides; absent/parse-fail -> all-
defaults, never fatal)
53 @backend/config/loader.py::load_config ->
@backend/config/models.py::AppConfig
54 -> video_id = MD5(whole file) @backend/utils/hashing.py::compute_video_id
(single shared helper, 64KB chunks; imported by main + all run_*.py ? no more per-entrypoint
copies)
55 -> validate
@backend/pipeline/validators/base.py::registry.run_all_with_context(video_path, global_config)
56 FormatValidator -> ctx['ffprobe_meta'] -> ResolutionFPSValidator -> PTSContinuity
-> SceneCut -> Blur -> Relevance (? ACTUAL registration/exec order ? see
@backend/pipeline/validators/__init__.py)
57 [validation FAIL -> persist status="failed", return early; report STILL emitted
in finally]
58 -> db.add_video(video_id, filepath, status, [r.dict()])
@backend/infrastructure/database.py
59 -> seg_name = req.segmenter or global_config.indexing.default_segmenter (fallback
'motion')
60 -> seg = segmenter_registry.get(seg_name)(db, global_config) ?
@backend/pipeline/segmenters/base.py::segmenter_registry (400 + available list if missing)
61 -> segments = seg.process_video(video_id, video_path, player_pool?, max_frames?)
62 [STOP POINT: segments-only ? predictions now in DB; eval/regression operate here
without stitching]
63 -> finally: emit_indexer_report(...) @backend/reporting/__init__.py (? ALWAYS runs

```

```

? even on validation failure / exception; NEVER raises; grafts response["reports"])
64     -> select segment ids -> [(start,end)] (+ stitching padding)
65     -> VideoCompiler.compile_highlights(filepath, time_pairs, out_name)
@backend/pipeline/stitcher/compiler.py
66     per-clip ffmpeg re-encode (libx264/superfast/crf22) -> concat -c copy ->
output/<name>.mp4
67     ```
68     Hybrid segmenter internal 4-phase (yolo_hybrid / tracknet_hybrid / wasb_hybrid ? same
shape):
69     ```
70     P1 chunk video (yolo only; tracknet/wasb = whole-video single pass, chunk_index=0)
71     -> P2 candidate "action windows" from motion/trajectory [STOP: candidates persisted
via db.add_candidate_segment]
72     -> P3 pad clip (ai_handoff_padding) -> ThreadPoolExecutor ->
provider.analyze_video(clip, prompt) ? provider_registry
73     -> P4 db.add_play_segment + db.link_candidate_to_segment
74     ```
75     Pure-AI path `gemini`: no P2; chunk-or-single -> provider -> heavy validate
(clamp/dedup) -> add_play_segment.
76     Legacy `motion`: pixel-diff -> segments + **fabricated** metadata/events (? not ground
truth).
77
78     ### 1B. WASB shuttle substrate (P2 of wasb_hybrid)
79     ```
80     WasbHybridSegmenter.process_video
@backend/pipeline/segmenters/wasb_hybrid.py
81     -> WasbRunner.healthcheck() (gate; not ok -> return [])
@backend/pipeline/detectors/wasb_runner.py
82     -> WasbRunner.run_predict(video, out_dir)
83     stage video + wasb_infer.py into WSL fs -> `wsl bash -c 'source conda.sh && python
wasb_infer.py ...`
84     -> @backend/pipeline/detectors/wasb_infer.py::run (WSL): sliding windows -> GPU
detector pass (CHECKPOINTED det_raw.jsonl) -> CPU tracker (always full re-run) -> trajectory CSV
85     -> TrackNetRunner.parse_trajectory_csv(csv)
@backend/pipeline/detectors/tracknet_runner.py (shared, re-exported by WasbRunner)
86     -> TrackNetRunner.trajectory_to_action_windows(points, fps, frame_width, chunk_start,
velocity_thresh, merge_gap, min_window_duration, chunk_index) [the model-agnostic WINDOWING
BRAIN]
87     ```
88     (`tracknet_hybrid` identical, swapping WasbRunner->TrackNetRunner; `_probe_fps_width`
fallback (30.0, 1280.0).)
89
90     ### 1C. Calibration / eval flow (NO pipeline run unless told)
91     ```
92     GT CSV -> annotations.load_annotations / calibrate_wasb.load_golden_gts ->
List[Interval]
93     DB predictions -> harness.get_predictions(db, video_id, stage) -> List[Interval] stage
? {candidates, final}
94     -> metrics.match_intervals (greedy temporal-IoU) -> metrics.score -> Score
95     -> harness.evaluate -> EvalReport -> report_to_dict
96     -> batch.aggregate (corpus micro/macro) @backend/eval/batch.py
97     -> regression.regress(_with_provider) vs baseline @backend/eval/regression.py
[CI gate]
98     Calibration: calibrate_wasb.run -> calibration.{select_best, improvement_report,
cross_validate} @backend/eval/calibration.py
99     Distill Gemini->local: calibrate_local (thresholds) OR distill_local (learned
classifier) @backend/eval/{calibrate_local,distill_local}.py
100    Learned segmenter (offline): training.build_training_set(X,y) ->
classifier.LogisticRegression.fit/save
101    Gemini refinement driver (tuning, standalone):
gemini_refine.{refine_window,full_video_rallies} @backend/eval/gemini_refine.py
102    Audio probe (diagnostic only, NOT in live pipeline): audio/e1_probe.run
@backend/eval/audio/e1_probe.py
103    ```
104    Entrypoints: `@run_eval.py` (single video score), `@scripts/run_phase3_eval.py`

```

```

(manifest batch, can segment), `@run_regression.py` (baseline gate, exit 1=regression),
`@backend/eval/calibrate_wasb.py` + `@backend/eval/export_wasb_segments.py` (WASB tuning
drivers).
105
106     ---
107
108     ## 2. SUBSYSTEMS
109
110     ### 2.0 Orchestration & config
111     - `@backend/main.py` ? Thin CLI + FastAPI bootstrap (orchestrator; modes via FLAGS).
    Static UI mounts `/static` from `backend/api/static`. Logging is initialized with `basicConfig`
    (INFO, timestamped format); backend modules use `logger = logging.getLogger(__name__)` (replaces
    stray prints in hot paths like motion, stitcher, wasb_infer, config loader). User-facing CLI
    summaries may still use `print` for direct output.
112     - `::app` FastAPI; `::db_instance`, `::global_config` (`Optional[AppConfig]`) module
    globals set ONLY in `main()`. ? importing `app` for ASGI without `main()` -> both `None` ->
    every endpoint NPES.
113     - **ASYNC (#16, PR-1):** `::process_video` `@POST /api/process` -> fast-fail 4xx
    (path/exists/segmenter) -> `db.create_job` -> `_get_executor().submit(_run_job,...)` -> **202 +
    job_id**. `::run_job` (worker: mark running -> `_execute_processing` -> done/failed; catches +
    logs traceback). `::execute_processing` = the former sync body (hash -> validate -> add_video
    -> segment -> `finally:` always `emit_indexer_report`, grafts `response["reports"]`; never
    raises from the report); a RAISING segmenter now prune_after-restores pre-run rows before re-
    raising. `::get_job` `@GET /api/jobs/{job_id}` -> `{status, progress, video_id, error, result,
    reports, ...}` ? job `status` = EXECUTION state (`queued|running|done|failed`); the pipeline
    verdict lives in `result["status"]`. `::get_executor` lazy module-global
    `ThreadPoolExecutor(max_workers=1)` (single box + Gemini serial-upload learning; tests
    monkeypatch it inline). `main()` runs `db.fail_stale_jobs()` at startup (orphaned queued/running
    -> failed). `::compile_highlights` `@POST /api/compile`. `::query_play_segments` `@POST
    /api/query`. `::get_config` `@GET /api/config`. `::run_cli_diagnose_clip` `--debug-clip`: lazy
    `import cv2`, samples ±3s vs `indexing.motion_threshold` (dual model/dict read). ? CLI `--video-
    path` mode still runs the pipeline synchronously inline (separate code path, unchanged by #16).
114     - **UPLOAD/DOWNLOAD (PR-2, B2+B3):** chunked upload ? `::upload_init` `@POST
    /api/upload/init` (`{filename, total_size_bytes}` -> ext allowlist + size gate ->
    `{upload_id}`, `::upload_part` `@PUT /api/upload/{id}/part/{n}` (raw-bytes body, STRICTLY
    sequential n, streams to `.partial-<id>` under `security.upload_dir`; per-part/total cap breach
    -> 413 + upload ABORTED), `::upload_complete` `@POST .../complete` (`{total_parts}` -> assemble
    -> never-overwrite rename -> `{path, video_id(MD5), size_bytes}` ? client then POSTs
    /api/process with that path), `::upload_abort` `@DELETE /api/upload/{id}`. ? upload state
    (`::uploads` + `::uploads_lock`) is IN-PROCESS ? restart aborts partials (client re-inits);
    per-user partitioning lands with PR-3. `::download_highlight` `@GET
    /api/highlights/{video_id}/{filename}` ? safe-name + `resolve_under_root(output_dir)` then
    FileResponse stream (the old /api/compile alert path was browser-unreachable). Config:
    `SecurityConfig.{upload_dir='output/uploads', max_upload_mb=4096, max_part_mb=95,
    allowed_upload_extensions=[mp4,mov]}` (95MB: free-tier edge proxies cap bodies ~100MB).
115     - `::load_config` ? ? legacy thin wrapper returning a **dict**
    (`load_app_config(path).model_dump(...)`); kept for transition ? new code imports
    `load_app_config`/`AppConfig` directly. `::ProcessRequest`/`::QueryRequest`/`::CompileRequest`
    pydantic bodies.
116     - `main()` sets `global_config.output.output_dir = args.output_dir` (typed field on
    `OutputConfig`); `db_path = os.path.join(global_config.output.output_dir, output.db_name)`.
    `compile_highlights` reads the SAME `global_config.output.output_dir`, so `--output-dir` is
    honored end-to-end (the old write-only `_runtime_overrides` hack is gone).
117     - ? `--segmenter` argparse `choices` frozen at build time from
    `segmenter_registry.list_available()` ? a lazily-registered segmenter wouldn't appear. CLI
    default_segmenter fallback `motion`.
118     - `@scripts/run_process_and_stitch.py` ? one-shot demo; ? hardcoded paths, interactive
    `input()` (blocks automation), skips validation. Config via `backend.config.load_config`
    (typed); segmenter, padding, and db path all read from the model (no literal fallbacks).
119     - `@run_eval.py::main` ? score one video's DB preds vs GT CSV.
    `--stage{candidates,final,both}`. Exit 2=arg/IO. ? pure scorer; errors if video never processed.
    `::default_db_path` resolves via `backend.config.load_config`
    (`output.output_dir`/`output.db_name`). `get_file_md5` is an alias for `compute_video_id`.
120     - `@scripts/run_phase3_eval.py::main` ? walk manifest, segment-if-needed, score,
    aggregate. `load_dotenv()` at import. ? DB key = file MD5, NOT manifest's `video_id` (kept as

```

```

label); config via `backend.config.load_config`; `segmenter_cls(db, cfg)` receives the typed
`AppConfig` (same object the live pipeline feeds segmenters).
121 - `@run_regression.py::main` ? `regress_with_provider` vs baseline. Exit codes semantic:
**0=ok, 1=REGRESSION (CI gate), 2=missing DB**. ? `--update-baseline` runs AFTER compare
(snapshots even a regressing run); `::default_db_path` resolves via
`backend.config.load_config`.
122
123 - `@backend/config/__init__.py` ? ? canonical config import surface. Re-exports
`load_config`, `save_default_config` (from `.loader`) + `AppConfig`, `Config`, `IndexingConfig`,
`StitchingConfig`, `ValidationConfig` (from `.models`); `__all__` = exactly those 7. Use `from
backend.config import load_config, AppConfig`.
124 - `@backend/config/models.py` ? Pydantic v2 models; **single source of truth for
defaults**.
125 - `::AppConfig` ? root. Scalars `sport='badminton'`, `ai_provider='gemini'`,
`ai_model='gemini-2.5-flash'`; sections `validation`, `indexing`, `output`, `stitching`
(`default_factory`). `model_config = {extra:'allow', validate_assignment:True}`. ?
`extra='allow'` -> unknown/typo'd config.json keys silently kept (legacy tolerance).
126 - ? `::AppConfig.get` ? legacy dict-compat shim: returns nested **sections as plain
dicts** (`model_dump(mode='json')`, NOT the model) so chained
`config.get("indexing", {}).get(...)` keeps working; falls back to searching a leaf key in any
section (sections discovered dynamically from `model_fields`). ? "bit every validator on real
runs" per docstring ? returning the model breaks `.get`. `::AppConfig.__getitem__` raises
`KeyError` for unknown top-level keys else delegates to `.get`.
127 - `::IndexingConfig` ? segmenter/detector knobs (`motion_threshold=0.08`,
`min_segment_duration=3.0`, `dead_time_merge_gap=2.0`, `chunk_duration=90.0`,
`chunk_overlap=10.0`, `yolo_velocity_threshold=0.15`, `yolo_frame_skip=3`,
`ai_handoff_padding=2.0`, `ai_max_workers=5`, `log/store_yolo_candidates=True`, nested
`tracknet`). ? `default_segmenter='yolo_hybrid'` HERE but config.json overrides to `gemini`. ?
`::IndexingConfig.segmenter_must_be_registered` `@field_validator` is a **no-op pass-through** ?
a typo'd segmenter validates fine, only fails later at `segmenter_registry.get(...)`.
128 - `::ValidationConfig` (`min_resolution_width=1280`, `min_resolution_height=720`, `min_fps=29.9`,
`min_blur_threshold=20.0`, `max_scene_cuts_per_minute=0.5`, nested `relevance`);
`::ValidationRelevanceConfig` (court HSV gates, `court_pixels_min_ratio=0.05`).
`::TrackNetConfig` (WSL invocation contract: `wsl_distro='Ubuntu'`, `repo_dir`,
`conda_env='TrackNetV4'`, `weights_path='', `queue_length=5`, `velocity_threshold=0.02`).
`::OutputConfig` (`default_highlights_name`, `db_name='sports_indexer.db'`,
`output_dir='output'` ? CLI `--output-dir` overrides it on the typed model in `main()`).
`::StitchingConfig`
(`begin_padding=0.5`, `end_padding=0.5`, `use_cache=False`, `min_shot_count=1`). `::Sport` Literal
of 5 sports. `::Config = AppConfig` alias.
129 - `@backend/config/loader.py` ? load/save with defaults-merge.
130 - `::DEFAULT_CONFIG_PATH = Path("config.json")` (? CWD-relative, NOT repo-root).
`::get_built_in_defaults` = `AppConfig().model_dump(...)`
131 - `::load_config` ? precedence built-in-defaults -> file overrides. ? **shallow top-
level merge** (`{**defaults, **file_data}`): a section in config.json REPLACES the whole
defaults dict for that section (Pydantic then re-applies field defaults for missing leaves).
Missing/unreadable/parse-error file -> fully-defaulted AppConfig + `[CONFIG WARNING]` (non-
fatal).
132 - `::save_default_config` ? writes fresh `config.json`, **overwrites** if present
(used by scripts/doctor.py/--setup). `::get_default_config` ? ? transitional plain-dict alias.
133 - `@config.json` ? on-disk config; mirrors `AppConfig`. ?
**`indexing.default_segmenter='gemini'`** diverges from model default `yolo_hybrid` ? file
wins at runtime.
134 - `@scripts/doctor.py` (formerly root `setup.py`) ? `::init_default_config` ? delegates
to `backend.config.save_default_config` (single-source defaults; skips if config.json exists).
`::run_pip_install` GPU-aware (`nvidia-smi` -> cu124 else cpu). `::run_diagnostics` (Python>3.8,
ffmpeg/ffprobe on PATH, optional torch+CUDA). Run via `python scripts/doctor.py` or `indexer
--setup`.
135
136 ### 2.1 Segmenters (under pipeline/segmenters) ?
137 - `@backend/pipeline/segmenters/base.py`
138 - `::BasePlaySegmenter` ABC ? `__init__(db, config)`; abstract
`name`, `description`, `process_video`. $ `process_video(video_id, video_path, player_pool=None,
max_frames=None)` -> List[dict{id,start_time,end_time,duration,metadata}]`.
139 - `::SegmenterRegistry`, `::segmenter_registry` (singleton). ? **register a new

```

segmenter**: subclass `BasePlaySegmenter`, call `segmenter\_registry.register(MyCls)` at module bottom, AND import the module in `@backend/pipeline/segmenters/\_\_init\_\_.py` (registration is an IMPORT side-effect). ? `register()` builds `cls(None, {})` just to read `.name` -> name/description/`\_\_init\_\_` MUST NOT touch `self.db`/`self.config`. ? `list\_available()` re-instantiates + swallows exceptions into "No description available." (a throwing `description` is masked).

140 - `@backend/pipeline/segmenters/\_\_init\_\_.py` ? imports all 6
(`motion,gemini,yolo\_hybrid,tracknet\_hybrid,wasb\_hybrid,fusion\_hybrid`) = the de-facto registration manifest; re-exports them + `segmenter\_registry` via `\_\_all\_\_`.

141 - `@backend/pipeline/segmenters/trajectory\_hybrid.py::TrajectoryHybridSegmenter` ? SHARED ENGINE for the trajectory hybrids (2026-06-11 collapse of the copy-paste twins; the "fix BOTH" footgun is retired): healthcheck gate -> P2 trajectory->windows (+persist) -> P3 provider handoff -> P4 DB/link; single `:::shot\_count\_from` (old wasb/tracknet shot_count drift gone). NOT registered itself. ? Subclass contract: `family` (config section `indexing.<family>.velocity\_threshold(0.02)`, `output/<family>` trajectory dir, `<family>-hybrid-<model>` metadata tag), `display\_label` (log lines), `name`/`description` properties, `\_build\_runner(idx\_cfg)` -> (DetectorRunner, detection_params extras). **Two fusion seams (IDENTITY by default ? twins byte-identical, pinned by the equivalence test): `:::enrich\_candidate\_stats(cw,points,fps,frame\_width,video\_path,idx\_cfg)`** (base = the 4 motion-key stats dict; computed once into `window\_stats`, used for persist + the handoff gate) **and `:::select\_for\_handoff(windows>window\_stats,idx\_cfg)`** (base = all indices). ? shared knobs still YOLO-named (`log\_yolo\_candidates/store\_yolo\_candidates`); whole-video single pass; needs `{PROVIDER}\_API\_KEY` (missing -> []). Equivalence pinned by `tests/test\_trajectory\_hybrid\_equivalence.py` (detector identity, shot_count, offsets, + the base-hook-identity tests ? parametrized over both twins).

142 - `@backend/pipeline/segmenters/fusion\_hybrid.py::FusionSegmenter` (name `fusion\_hybrid`, **default-OFF** multi-signal fusion P2) ? subclasses the engine; `family='fusion'`, `display\_label='Fusion'`. `:::build\_runner` selects the substrate via `indexing.fusion.substrate` ('wasb'|'tracknet') reusing the existing runner+config. `:::enrich\_candidate\_stats` overrides to ALWAYS add `net\_crossings` (Gate A signal, GPU-free via `rally\_gate.gate\_windows(..., estimate=self.axis\_estimate(points))` ? axis/net cached per trajectory) and, only when `fusion.cue\_flags.person`, the ~5 person features (`:::person\_features` ? lazy `PersonDetector.detections\_per\_frame` over the window's ORIGINAL-video frame range ? `fusion\_features.compute\_person\_features`). `:::select\_for\_handoff` overrides to apply served Gate A (`fusion.min\_crossings`) ? Gate B (`fusion.min\_players`, person-only) ? default 0/0 ? all windows kept. ? persisted candidates are ALWAYS the full superset (offline substrate intact); only the handoff/play_segments are gated. ? to A/B vs a tuned wasb run, set `fusion.velocity\_threshold = wasb.velocity\_threshold` (the engine reads `idx\_cfg['fusion'].velocity\_threshold`). ? person features eager even for gated windows (P3 cleanup). Tests: `tests/test\_fusion\_hybrid.py`.

143 - `@backend/pipeline/segmenters/wasb\_hybrid.py::WasbHybridSegmenter` (name `wasb\_hybrid`, the LIVE shuttle substrate; MIT, ships pretrained badminton weights) ? thin subclass: `family='wasb'`, wires `WasbRunner`/`WasbConfig`; `detection\_params` extras = `weights\_path` only.

144 - `@backend/pipeline/segmenters/tracknet\_hybrid.py::TrackNetHybridSegmenter` (name `tracknet\_hybrid`) ? thin subclass: `family='tracknet'`, wires `TrackNetRunner`/`TrackNetConfig`; extras = `weights\_path` + `queue\_length`.

145 - `@backend/pipeline/segmenters/yolo\_hybrid.py::HybridYoloSegmenter` (name `yolo\_hybrid`, no YOLO remains; ADR-005). **Stage-1 detection EXTRACTED to `detectors/person\_detector.py` (fusion P1, no behavior change)** ? the segmenter now ORCHESTRATES only (chunking, pipelined-vs-sequential P2 prefetch where ffmpeg overlaps sequential GPU, AI-provider handoff, candidate persist/link) and delegates detection to `self.person` (`PersonDetector`). Imports `\_DEFAULT\_DETECTOR` from person_detector for the `detection\_params` snapshot. ? `yolo\_tracker`/`bytetrack.yaml` is dead/back-compat; `gemini\_` config keys are deprecated aliases still honored; still mpy-quarantined for pre-existing orchestration debt (the torchvision code that caused it left).

146 - `@backend/pipeline/segmenters/gemini.py::GeminiPlaySegmenter` (name `gemini`) ? pure-AI, no P2. Chunks are RE-ENCODED + downscaled to `indexing.ai\_chunk\_scale\_width` (default 1280; 0 = legacy stream-copy) before upload ? stream-copied 4K chunks exceed the provider's `MAX\_UPLOAD\_MB` and every chunk gets refused ("0 segments / success"); re-encode also fixes keyframe drift at chunk boundaries. Oversize refusals (provider metrics `oversize\_skipped`) are aggregated -> `logger.error` + `reporting.run\_signals` (`oversize\_clips\_skipped`) for the report rule. Workers from `indexing.ai\_max\_workers`. ? chunks only if `duration>chunk\_duration`; `max\_segment\_duration=90.0` hardcoded; `debug\_duration` silently truncates; `parse\_time` accepts colon timestamps (warns); does NOT persist candidate segments (only hybrids feed the fine-tuning

```

table).
147     - `@backend/pipeline/segmenters/motion.py::MotionPlaySegmenter` (name `motion`) ?
pixel absdiff baseline. ? **metadata + serve/smash/foul/winner events are RANDOM**
(random.sample/randint/uniform/choice); only segmenter using `db.add_event` and passing
server/receiver; uses `print()`; honors `max_frames`.
148
149     ### 2.2 Detectors & windowing (WSL-quarantined; ABC abstraction complete)
150     - `@backend/pipeline/detectors/base.py` ? DetectorRunner ABC (healthcheck() ->
Tuple[bool,str] never-raises; run_predict(...) -> Optional[str] CSV path). Contract deliberately
matches the pre-existing tuple returns and call sites in the hybrids. Rich module docstring
includes "how to add a new runner", Linux-native path notes (the de-risk goal), and guidance on
threading health into reports. Re-exported from `detectors/__init__.py`. This finishes the WSL
de-risk seam (low-regret 6 / review item 1).
151     - `@backend/pipeline/detectors/tracknet_runner.py`

```

9. user (2026-06-16T03:12:21.009Z)

```

1     # Multi-signal rally fusion ? add a human (person) cue alongside the shuttle
2
3     **Status:** **P1 DONE** (2026-06-13) ? the person cue is extracted to
4     `backend/pipeline/detectors/person_detector.py` (`PersonDetector`), consumed by
`yolo_hybrid` with no
5     behavior change. **P2?P4 remain owner-gated.** Phases below are sequenced but each is a
separate PR.
6     **Owner decisions captured** in [S9](#9-open-decisions-for-the-owner).
7
8     > **Companion docs:** the 2-layer mental model is
[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md);
9     > the modular-fusion architecture this builds on is **ADR-007**
([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md));
10    > the permissive-dependency rule is **ADR-002 / ADR-005** ; the audio cue (the #1 next
cue, already
11    > scaffolded) is [archives/research/AUDIO_RALLY_DETECTION_PLAN.md](archives/research/AUD
IO_RALLY_DETECTION_PLAN.md).
12    > The vendor labels that gate the learned path come from
[GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md)
13    > and the labeling spec in [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md).
14    > **How each cue here is added/tested without production regression** ? A/B, shadow
eval, and the
15    > promote-only-on-lift discipline ? is [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)
(every P1?P4 cue
16    > below lands flag-gated, default OFF, promoted only on measured LOVO lift).
17
18    ---
19
20    ## 1. Context ? why this, why now
21
22    Rally detection today runs **one cue at a time**. Each segmenter independently turns a
single signal
23    into rally windows:
24
25    | Segmenter | Signal | File |
26    |---|---|---|
27    | `wasb_hybrid` / `tracknet_hybrid` | shuttle velocity (frozen WASB/TrackNet trajectory)
| `backend/pipeline/detectors/tracknet_runner.py` |
28    | `yolo_hybrid` (the Pydantic `default_segmenter` fallback; `config.json` overrides it
to `gemini`) | person velocity (torchvision + ByteTrack) |
`backend/pipeline/segmenters/yolo_hybrid.py` |
29    | `motion` | raw pixel motion | `backend/pipeline/segmenters/motion.py` |
30    | `gemini` | a cloud VLM watching the whole clip | `backend/providers/gemini.py` |
31
32    **There is no fusion.** The owner's directive: make the detector **use the shuttle *and*
the humans
33    together** ? and be ready to fold in more signals (audio, later pose) ? delivered
34    **MIT / Apache-2.0 / BSD-clean** so we can ramp commercially without a licensing

```

```

rethink.
35
36     This is **not greenfield** ? it operationalizes work the repo already blessed:
37
38     - **`HOW_RALLY_DETECTION_WORKS.md` open question:** **Fusion: combine shuttle-motion
(WASB) with
39     player-motion for static cams."**
40     - **ADR-007** already designed a **modular "rally layer"*** that fuses cues *outside* the
frozen WASB
41     model, with **audio as the #1 next cue and person/pose parked as #2**. This plan picks
up the parked
42     person cue and generalizes the fusion seam so future cues are "register a producer,"
not "re-touch
43     the segmenter."
44     - **`backend/pipeline/detectors/rally_gate.py` already exists** ? the net-crossing gate,
explicitly
45     labelled in its own docstring as **Gate A in the over-fire fix spectrum (**B = player-
stance state
46     machine, future**).** The person cue in this plan **is that named-but-unbuilt Gate
B.**
47
48     ### The owner's "held-shuttle" false positive ? what's actually going on
49
50     The owner observed a "rally" that was just a player **holding / flicking the shuttle in
hand** between
51     points. This is the *exact* failure `rally_gate.py` was written to kill:
52
53     > **WASB over-detects: when a player flicks/tosses the shuttle back to the opponent for
service after
54     > losing a point, that single trajectory looks like a rally. The discriminator is
structural: a real
55     > rally has the shuttle crossing the net MULTIPLE times; a single service feed crosses
~once."**
56
57     **The structural fix already exists and is validated offline ? but it is NOT wired into
the production
58     runtime path.** `rally_gate.apply_gate(..., min_crossings=2)` is called only from the
**eval drivers**
59     (`eval_partial_labels.py`, `export_wasb_segments.py`, `calibrate_local.py`;
`distill_local.py` uses
60     `rally_gate.gate_windows` to compute the `net_crossings` feature, not `apply_gate`); the
live segmenters
61     historically didn't apply it. P2 closed this: `FusionConfig.min_crossings` now exists
(default `0` =
62     OFF), so the served Gate A is a config flag ? consistent with $10 P2 below, marked DONE.
So the
63     single cheapest, highest-confidence win in this whole plan is **P2's first step: fold
the existing
64     Gate A into production.** Everything else (the person cue / Gate B, the learned fusion)
is robustness
65     on top of that.
66
67     ### Why humans help ? and the honest risk
68
69     The beachhead videos are **static tripod** (Bangalore academies,
[PMF_MARKET_RESEARCH.md](PMF_MARKET_RESEARCH.md)),
70     so player presence/motion is a usable cue. ADR-001's warning was specifically about
**broadcast pan**
71     cameras, where player-motion fooled the old heuristic ? *that failure mode does not
apply to a fixed
72     court cam.* But we keep **the shuttle as the camera-invariant primary** so the general
bet stays robust
73     on other camera types. Humans then add:
74
75     - **Precision** ? reject shuttle-on-floor / warm-up knock-ups / held-shuttle feeds when

```



```

the court
76     isn't occupied by ?2 players in a play posture.
77     - **Recall** (via the learned classifier) ? features that help bridge shuttle
occlusion/blur gaps.
78
79     ---
80
81     ## 2. Core architecture ? one rally-fusion layer, cues are swappable producers
82
83     Formalize ADR-007's diagram. Every cue is a **black box that emits a time-aligned
signal** (per-frame
84     activity and/or per-window features); **fusion happens in *our* layer, never inside a
frozen model.**
85
86     ```
87     frames ?? WASB (frozen, MIT)      ?? shuttle trajectory ???
88     frames ?? person detector (ours)  ?? tracks + boxes ??????
89     audio ?? hit detector (ours)      ?? hit events ???????????? RALLY LAYER (ours) ??
rallies
90     [frames ?? pose/stance (ours; parked #3) ?? stance ???????? feature + decision fusion
91     ```
92
93     - **Shuttle = primary, camera-invariant.** Windows still **seed** from
94     `TrackNetRunner.trajectory_to_action_windows()` (`tracknet_runner.py:234`). Other cues
*enrich and
95     adjudicate* shuttle-seeded windows; they do **not** replace the shuttle seed. Person
runs only over
96     **shuttle-seeded candidate windows + padding**, not the whole video ? which bounds the
added cost.
97     - **A small `RallyCue` contract.** Each producer yields `(frame_idx ? features)` + per-
window stats.
98     Adding a cue = register a producer; the existing `SegmenterRegistry`
(`segmenters/base.py:35`) is the
99     template for the pattern. Keep the single-cue segmenters registered as fallbacks.
100
101     ---
102
103     ## 3. The two fusion modes (both ? per the owner decision)
104
105     ### 3.1 Decision-level gate (safe baseline, ships first, no labels needed)
106
107     ```
108     rally_active = shuttle_motion ? (?2 persons in court zone ? recent person track)
109     ```
110
111     The `? recent-track` clause tolerates 1?2 frame person-detector misses so the AND-gate
can't cause a
112     recall cliff. This composes with the **existing net-crossing Gate A**:
113
114     ```
115     keep window ? net_crossings ? 2                (Gate A ? EXISTS in rally_gate.py, eval-
only today)
116     ? (?2 in-court players ? recent track)    (Gate B ? the person cue, NEW)
117     ```
118
119     Gate A kills the service-feed / held-shuttle case structurally; Gate B adds court-
occupancy evidence
120     for the cases a single trajectory can still fool (e.g. a knock-up that does cross the
net). No training
121     data required ? this is the label-free fusion that ships before any golden labels exist.
122
123     ### 3.2 Feature-level (primary, the learned path)
124
125     The cues emit a **small set of aggregated, scale/translation-invariant per-window
features ? NOT raw

```

```

126     coordinates.** Raw pixel coordinates would memorize *this* court/camera and fail to
generalize off a
127     still-small labeled set; counts / ratios / zone-membership / two-sidedness transfer.
128
129     **What already feeds the classifier** (`distill_local.py:40`, `FEATURE_NAMES`):
130     `duration, peak_velocity, mean_velocity, active_ratio, net_crossings` ? note
**`net_crossings` is
131     already a feature** (computed via `rally_gate.gate_windows`).
132
133     **Person features to add (~4, deliberately few):**
134
135     | Feature | Meaning | Rally signal |
136     |---|---|---|
137     | `players_in_court` | median count of persistent in-court tracks | ? 2 (singles) / 4
(doubles); 0?1 = dead time |
138     | `two_sided_motion` | are moving players on **both** net-halves? | rally = two-sided;
one person feeding = one-sided |
139     | `mean_player_motion` | aggregate normalized player velocity | rally = sustained;
resting/standing = low |
140     | `in_court_ratio` | fraction of frames with ?2 in-court players | track stability vs
flicker |
141     | `shuttle_player_colocation` | fraction of frames the shuttle sits inside/adjacent to a
person box (hand region) | **held = high; in-flight = mostly free space** ? the direct held-
shuttle discriminator |
142
143     These join the shuttle features in the candidate `stats` JSON
144     (`database.py::add_candidate_segment`) and the feature vector of the **existing
distilled classifier ?
145     the tiny numpy-only L2 logistic regression in `backend/eval/classifier.py` (ADR-004),
NOT a new net.**
146     It scores each candidate window `P(real rally)` and replaces the Gemini call at runtime.
147
148     **What the model learns:** the **weights** on those few features ? e.g. down-weight a
shuttle-motion
149     window with 0?1 in-court players / one-sided motion / high shuttle-colocation (warm-up,
floor shuttle,
150     held-shuttle feed); up-weight ?2 players, two-sided, sustained motion, shuttle in free
flight.
151
152     **Small-N discipline (non-negotiable):**
153     - Keep to **~4 person features** ? more features on a still-small labeled set overfit;
that's exactly
154     why the model stays a logistic regression, not a net.
155     - Use the **leave-one-video-out (LOVO)** protocol that `distill_local.py` **already
implements**: train
156     on the other videos' candidates, predict the held-out *video*, score vs its labels ?
never a
157     held-out *window* from the same video (windows in one video are correlated; that
flatters the score).
158     - Training labels today are **Gemini silver labels** (distilled offline, ADR-004/007 ?
Gemini is a
159     teacher, never a runtime dependency, never a training-data *source* for owned
weights); **golden
160     labels are the honest measuring stick**, and as the Roora golden corpus grows they can
also become
161     training labels. Until the labeled set is larger, treat the learned weights as
162     **calibration / direction-check, not "a fully general model"** ? the decision-level
gate (§3.1) is
163     the shippable, label-free fusion in the meantime.
164
165     > *(Gated on golden labels existing ? the same blocker ADR-007 flags for audio. See
166     > [GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md).)*
167
168     ---
169

```

```

170    ## 4. Shuttle-state: in-flight vs in-hand (the misclassification, mapped to real code)
171
172    The "shuttle moved" seed conflates a shuttle **in flight** with a shuttle **carried in
hand**. The
173    discriminators, and where each lives:
174
175    | Signal | In-flight | Held / static | Status |
176    |---|---|---|---|
177    | `net_crossings` | ? 1 (usually several) | ~0 (held) / ~1 (single feed) | **EXISTS**
(`rally_gate.py`); eval-only ? wire to prod |
178    | `shuttle_player_colocation` | low (free space) | high (near a hand) | **NEW** (needs
the person cue) |
179    | `direction_reversals` | several (racket contacts) | ~0 | derivable from trajectory;
partial overlap with crossings |
180    | `peak shuttle speed` | high | walking speed | derivable from existing velocity |
181
182    So the in-flight-vs-held fix is **mostly Gate A (already built) + the new colocation
feature from the
183    person cue**. **No new label column is needed** ? the golden labels already start a
rally at the
184    *serve-contact frame* and treat all pre-serve / held-shuttle time as dead time (see
185    [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md)), so held-shuttle moments
are negatives
186    by construction and are the ground truth that proves the detector wrong.
187
188    ---
189
190    ## 5. Adversarial mitigations (integrating safely)
191
192    - **Court-region mask** ? polygon around the court; drop persons outside it (spectators,
coaches,
193    adjacent courts). Config-supplied per camera; ~1 ms. Source: a manual per-camera
polygon first;
194    the **4 court corners + net-line** that Roora labels once per video (see the labeling
doc) build
195    this mask **and** the net split that `two_sided_motion` / `net_crossings` need.
196    - **Player-count constraint** ? expect **2 (singles) / 4 (doubles)** persistent in-court
tracks;
197    suppress transient/extra detections. The `singles_or_doubles` manifest flag sets the
expectation.
198    - **Temporal association** ? reuse `supervision` **ByteTrack** (already in
`yolo_hybrid`); reject
199    singleton detections; smooth across frames.
200    - **Multi-court rejection** ? proximity to the primary court region (homography
optional, later).
201    - **Domain shift (COCO frontal ? side/overhead court)** ? start **box-only** (robust);
keep the
202    detector swappable. Fine-tuning a person detector is **parked** ? we have no in-house
person labels,
203    and the **paid-LLM/training guardrail + minors-exclusion** apply to any future person-
training data.
204    - **Fusion-degradation guard** ? shuttle stays primary; person is gate + feature, never
the sole
205    trigger; the single-cue segmenters remain registered fallbacks.
206
207    ---
208
209    ## 6. Box vs pose
210
211    **Box-first.** Occupancy + count + two-sided motion + shuttle-colocation is enough for
both the gate
212    and the features, and box detection is the cheap, robust win. **2D pose is parked behind
the same
213    `RallyCue` interface ? pose (MediaPipe Pose / RTMPose, both Apache-2.0) buys serve-
ready stance

```

214 gating and shot cues later, at the cost of a second model. This matches ADR-007 parking
pose as the
215 #2/#3 cue.
216
217 ---
218
219 ## 7. Performance
220
221 The person detector is the **existing in-process torchvision model** run at **frame-skip**
375
222 (`yolo_frame_skip` default 3`) with track interpolation between sampled frames. The
shuttle keeps
223 running in its separate WSL process ? the two are **separate processes**, so a "shared
backbone" isn't
224 free; budget realistically as the existing `yolo_hybrid` cost minus skip savings. Person`
runs **only**
225 over shuttle-seeded windows + `ai_handoff_padding``, not the whole video, which bounds
the cost.
226
227 ---
228
229 ## 8. Licensing (the guardrail table ? all green)
230
231 Every dependency stays **MIT / Apache-2.0 / BSD** (ADR-002 / ADR-005; `ultralitics`` was
dropped for
232 AGPL-3.0).
233
234 | Cue / model | Code | Weights / data | Verdict |
235 |---|---|---|---|
236 | Shuttle: WASB-SBDT, TrackNetV3 | MIT | frozen badminton weights | ? (ADR-001/002) |
237 | Person (default): torchvision Faster R-CNN / RetinaNet / FCOS | BSD/MIT | COCO-
pretrained | ? already in repo (`_DETECTOR_FACTORIES``) |
238 | Person (escalation): **YOLOX**, **RT-DETR / RT-DETRV2** | Apache-2.0 | COCO /
Objects365 | ? |
239 | Pose (parked): **MediaPipe Pose**, **RTMPose / MMPose** | Apache-2.0 | COCO-structure
| ? |
240 | Tracking: `supervision` ByteTrack` | MIT | n/a | ? already in repo |
241 | **Banned** | **Ultralytics YOLOv5/8/11 (AGPL-3.0)** . **MonoTrack (Adobe, non-comm)**
· **YOLO-NAS (Deci, non-comm)** | ? | ? |
242
243 **COCO caveat to record:** COCO **annotations** are **CC BY 4.0** (commercial OK **with**
attribution);
244 COCO **images** are Flickr-ToS (per-image). Pretrained-**weight** use is clean; record the
attribution.
245
246 ---
247
248 ## 9. Open decisions (for the owner)
249
250 1. **Court-region mask source** ? manual per-camera polygon (simple, ship first) vs auto
court
251 homography/line detection (later). **Recommend manual first**, fed by Roora's court-
corner labels.
252 2. **Person detector default** ? keep torchvision (zero new deps, BSD/MIT) vs adopt
YOLOX/RT-DETR
253 (Apache, faster) now. **Recommend torchvision first**, escalate only if eval shows a
speed/accuracy gap.
254 3. **Pose now or parked** ? **recommend parked** behind the interface (box-first).
255 4. **Golden labels** ? feature-level fusion (P3) is **blocked on golden rally labels**;
the source is
256 the VLC `rally-annotator` + the Roora vendor pilot`
(`[GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md)`).
257
258 ---
259

```

260     ## 10. Phased execution (owner-gated; ONE PR per slice, off `master`)
261
262     - **P0 ? this design doc** +
[ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md), indexed in
263     `docs/README.md`, cross-linked from `HOW_RALLY_DETECTION_WORKS.md` and ADR-007. *(this
deliverable;
264     docs-only, no code.)*
265     - **P1 ? extract the shared person cue:** ? **DONE (2026-06-13).** Lifted the
torchvision detector +
266     ByteTrack + velocity out of `yolo_hybrid` into
`backend/pipeline/detectors/person_detector.py`
267     (`PersonDetector`); `yolo_hybrid` now orchestrates only and delegates to `self.person`
(no behavior
268     change ? regression-safe, proven by the unchanged process_video output tests). New
269     `tests/test_person_detector.py`; suite +1 green; person_detector is type-clean (left
mypy quarantine
270     for the moved code). The fusion segmenter (P2) reuses this same producer over shuttle-
seeded windows.
271     - **P2 ? the FusionSegmenter (Gate A + Gate B + person-feature persistence):** ? **DONE
(2026-06-14).**
272     Shipped `backend/pipeline/segmenters/fusion_hybrid.py::FusionSegmenter` (registry
`fusion_hybrid`,
273     **default-OFF** ? `default_segmenter` unchanged), subclassing the trajectory engine
via two identity
274     hooks (`_enrich_candidate_stats`, `_select_for_handoff` ? wasb/tracknet stay byte-
identical, adversarially
275     verified). It seeds windows from the shuttle substrate, **always persists
`net_crossings`** (Gate A signal),
276     and ? when `indexing.fusion.cue_flags.person` is on ? runs the person cue over each
window's frames
277     (`PersonDetector.detections_per_frame`) and persists the ~5 person features
(`fusion_features.py`,
278     pure + unit-tested). Served gating is optional + default-OFF: `fusion.min_crossings`
(Gate A) /
279     `fusion.min_players` (Gate B) skip sub-threshold windows from the handoff while
**persisting the full
280     superset** (offline substrate intact). Court mask = optional `fusion.court_polygon`
(None ? count-only).
281     Lift to be measured via the harness; held-shuttle-drop + person-feature values are
GPU/owner-side.
282     - **P3 ? feature-level (learned) fusion:** add the person features (incl.
`shuttle_player_colocation`)
283     to `distill_local` / `classifier`; recalibrate IoU on the golden set; report **LOVO
lift on held-out
284     videos**. *(Gated on golden labels.)*
285     - **P4 ? additional cues:** wire **audio** (ADR-007's #1, already scaffolded in
286     `backend/pipeline/audio/hit_detector.py`) and later **pose**, through the same
`RallyCue` interface.
287     Each is add-a-producer, not a refactor.
288
289     ---
290
291     ## 11. Verification
292
293     - **Offline, runnable in this dev container (no GPU/torch/cv2):** unit tests for the
fusion gate logic
294     and `RallyCue` aggregation (court-mask filter, player-count constraint, `? recent-
track` tolerance)
295     and the shuttle-state signals (net-crossing count ? extends the existing `rally_gate`
tests ? and
296     person-colocation) with **synthetic trajectories + person boxes**. A key-license
assertion test
297     (no `ultralytics` import; detector factories stay torchvision/Apache).
298     - **Machine-side (owner, GPU/WSL ? the full pipeline can't run in this container):** run
the

```

```

299     `FusionSegmenter` on a real static-tripod clip; confirm **held-shuttle / warm-up /
spectator false
300     positives drop** vs shuttle-only (this should reproduce the owner's observed bug and
show it fixed ?
301     start by confirming Gate-A-in-production alone fixes it); **LOVO IoU lift on the
golden set** for the
302     learned path.
303     - **Docs (P0):** links resolve; both docs indexed in `docs/README.md`; this doc cross-
linked from
304     `HOW_RALLY_DETECTION_WORKS.md` (the fusion open-question) and ADR-007.
305

```

10. user (2026-06-16T03:12:21.370Z)

```

1      # How rally detection works (the 2-layer mental model)
2
3      A short conceptual map of how we go from raw video to rally `[start, end]` segments ?
4      and, crucially, **which layer our annotation can improve.**
5
6      ...
7      video ??? frames ??? [DETECTION] ??? shuttle (x,y,visible) per frame
8              ?
9              (WASB CNN,          per-frame trajectory
10             frozen weights)      ?
11                                 ?
12                                 velocity signal (how fast the shuttle moves)
13                                 ?
14                                 [SEGMENTATION] ??? the 3 knobs we tune
15                                 ?
16                                 ?
17                                 rally windows [start, end] ??? IoU vs golden
labels
18      ...
19
20      ## Two layers ? keep them separate in your head
21
22      **1. DETECTION ? "where is the shuttle, this frame?"**
23      WASB is a CNN that, for each frame, outputs a heatmap ? the shuttle's `(x, y)` + a
24      visibility/confidence. Its weights were learned (by gradient descent, on shuttle-
coordinate
25      labels) and we use them **frozen**. We chose a *shuttle* detector on purpose: "shuttle
in
26      flight" is the signal for "a rally is happening," and it's **camera-agnostic** ? unlike
27      player-motion, which fooled the old heuristic on broadcast cameras that pan with the
players.
28      (Your videos are static tripod, so player-motion would also work ? but shuttle-motion is
the
29      robust general bet.)
30
31      **2. SEGMENTATION ? "when is a rally?"**
32      From the per-frame `(x, y)` we compute the shuttle's **velocity** (displacement between
frames).
33      Then three knobs turn that into rally intervals:
34      - **`velocity_thresh`** ? a frame is "active" only if the shuttle is moving faster than
this
35      (in flight, not lying on the floor / being picked up).
36      - **`merge_gap`** ? active frames within this many seconds get joined into one rally
(bridges
37      brief gaps: a shuttle momentarily occluded, or the beat between two hits).
38      - **`min_window_duration`** ? throw away windows shorter than this (blips / false
flickers).
39
40      Output: a list of `[start, end]` windows = the rallies.
41
42      ## Where the annotation fits (the key point)

```

```

43     Our Roora/golden labels are **rally `[start, end]` times** ? they say *when* rallies
are, not
44     *where the shuttle is*. So they can't teach WASB's detector (that needs `(x,y)` labels).
They
45     **can** grade and tune the **segmentation layer**: we score predicted windows against
the golden
46     rallies with **temporal IoU** (how much they overlap), and pick the 3 knobs that
maximise it ?
47     **calibrate on some rallies, measure on held-out ones** so the number is honest.
48
49     ## Where sparks could land (open questions ? not commitments)
50     Things worth a shower-thought; any of these could beat 3 hand-tuned thresholds:
51     - **Richer features than raw velocity:** acceleration, trajectory *continuity* (a real
rally is a
52     smooth arc that survives occlusion), shuttle *height/arc*, or **direction reversals**
(each
53     reversal ? a shot ? ties to `shots_count`!).
54     - **A tiny learned segmenter** instead of 3 thresholds: feed the trajectory features to
a small
55     model trained on the golden rally labels (still the segmentation layer, just smarter).
56     - **Hysteresis** (two thresholds: enter-rally high, stay-in-rally low) ? robust to mid-
rally dips.
57     - **Rhythm/periodicity:** rallies have a hit cadence; dead time doesn't.
58     - **Fusion:** combine shuttle-motion (WASB) with player-motion for static cams. ?
designed in
59     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (one rally-fusion layer;
the person cue is
60     ADR-007's parked #2 and `rally_gate.py`'s named-but-unbuilt "Gate B").
61     - **Per-video calibration as a *feature*:** "watch a few labeled rallies ? auto-tune to
this
62     court/lighting ? segment the rest sharply" ? exactly your earlier idea, at the right
layer.
63
64     The whole point of per-video calibration: a handful of *your* labeled rallies nudges
these knobs
65     to *your* court/camera, and we report the IoU lift on the rallies it never saw.
66

```

11. user (2026-06-16T03:12:21.499Z)

```

1     # A signal-fusion framework for rally detection ? analyzers, scopes & reconcilers
2
3     **Status:** **DESIGN ? owner-gated, NOT started.** A peer to
4     [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) and
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md).
5     This doc specs a small **framework** for evolving rally-detection quality: producers
that emit **signals**
6     at three temporal **scopes** (frame / window / session), a **scorer** that *learns to
weight* them, and an
7     auditable **reconciliation** pass that lints the decision against the evidence. The goal
is to keep piling
8     on clever cues **without the pipeline becoming a pile of special cases.** Docs-only; no
implementation
9     code. Each phase below is a separate, owner-gated PR.
10
11     > **Companion docs:** the detection-layer producers (shuttle / person / audio / shuttle-
state) live in
12     > [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (`RallyCue` = swappable
producer; the
13     > `_enrich_candidate_stats` hook; scale/translation-invariant feature discipline, §3.2);
the no-regression
14     > A/B + LOVO measurement these ride is [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)
(`Variant`, `compare`,
15     > `compare_unlabeled`); the scorer is the distilled logistic regression (**ADR-004**,
16     > `backend/eval/classifier.py`); the modular rally layer is **ADR-007**

```

```

17      > ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the 2-layer
mental model is
18      > [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md); the live quality log is
19      > [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md).
20
21      > **Framing.** Rally detection is a genuinely hard temporal problem ? weak, multi-modal
cues; interval-level
22      > ground truth that is scarce and expensive; a confound (warm-up / knock-ups / multi-
court / held-shuttle)
23      > that looks identical to real play in any single frame or frequency band. This doc is
written as a
24      > **reusable framework**, not a badminton-specific patch: the contracts are sport-
agnostic and the design
25      > is meant to generalize to other weakly-supervised temporal-segmentation problems
($10). Where that
26      > ambition touches research claims, $10 says plainly what is and isn't yet
substantiated.
27
28      ---
29
30      ## 1. The spine ? NO special-casing in the decision path
31
32      This is the one principle everything else serves:
33
34      > **Every producer (analyzer, detector, reconciler) emits a SIGNAL carrying a value + a
confidence/presence.
35      > The SCORING MODEL consumes those signals and *learns to weight them*. We do NOT add
hand-coded `if/else`
36      > branches into the decision.**
37
38      **Producers propose; the scorer arbitrates; the experiment harness + golden set decide
what gets promoted.**
39      That separation is what lets us add a net-hit service fault, a let-cord, a double-
bounce, a shuttle-out
40      cue, a *warm-up span*, a *shuttle-state* ? each as **register a producer + a feature
column,** never **add
41      a branch.** A signal that doesn't earn its keep stays at weight ? 0; it costs nothing
because it is an
42      unused column, not a live `if`.
43
44      The pattern already exists in miniature: `experiment.predict()`
(`backend/eval/experiment.py:213`) treats
45      `net_crossings` and `players_in_court` as **feature columns** ? `_gate_mask` is the
trivial scorer
46      (threshold a column), `LogisticRegression` (ADR-004) the richer learned scorer (weight
all columns), a
47      future scorer richer still. The decision path is *signals ? scorer*, with **zero** `if
service_fault`:
48      drop.. Analyzers feed it more columns; reconcilers run *outside* it as an audited post-
pass.
49
50      ...
51      DETECTION ? frame-scope producers (expensive, GPU/WSL, run ONCE)
52      WASB shuttle (x,y,v,vis) · person boxes+tracks · audio hits · shuttle-
state(in_air/contact/inactive)*
53      ? per-frame-aligned signals (+ net estimate, window
bounds)
54      ?
55      shuttle-seeded candidate windows (FusionSegmenter, P2)
56      ?
57      ANALYZERS ? decision-layer producers (pure, offline, re-scorable) [forward
seam]
58      ?? window-scope (WindowAnalyzer): one candidate ? features + events + boundary
hints e.g. service-fault
59      ?? session-scope (SessionAnalyzer): whole timeline ? labeled spans + per-window

```



```

membership e.g. warm-up
60                                     ? every signal ? candidate_segments.stats as feature
columns
61                                     ?
62         SCORING / DECISION ? signals ? learned scorer (weights, NO branches)
63         _gate_mask (trivial) + LogisticRegression (ADR-004) + harness-pinned model
64                                     ? predicted rally segments
65                                     ?
66         RECONCILERS ? backward seam, "linter for rally decisions" (report-first, idempotent)
67         ConsistencyRule registry: decision-vs-evidence conflicts ? Issue + corrected set
68                                     ?
69         stats · events · spans · corrections record · highlights      EXPERIMENT HARNESS
gates EVERY promotion
70         (* shuttle-state = a richer frame detector; its PRODUCER lives in
MULTI_SIGNAL_FUSION_PLAN, consumed here ? §3.5)
71         ```
72
73         ---
74
75         ## 2. The framework ? signals, scopes, and who arbitrates
76
77         The architecture has exactly three kinds of moving part. Capability is added by
registering a producer;
78         the decision logic never grows a branch.
79
80         **Producers emit signals at three temporal SCOPES.** A signal at a coarser scope is
computed by aggregating
81         finer-scope signals; everything ultimately lands as a per-candidate feature column (or
an event/span the
82         scorer or a human consumes).
83
84         | Scope | Producer | Sees | Emits | Cost | Examples |
85         |---|---|---|---|---|---|
86         | **frame** | detector (`RallyCue`) | raw frames / audio | per-frame signal |
**expensive** (GPU/WSL) | WASB shuttle, person boxes, audio hits, **shuttle-state** (§3.5) |
87         | **window** | `WindowAnalyzer` (§3.1) | one candidate window's aligned frame-signals |
features (+conf), events, boundary hints | cheap, pure | service-fault, double-bounce, let-cord,
shuttle-out |
88         | **session** | `SessionAnalyzer` (§3.4) | the whole timeline of windows | labeled spans
+ per-window membership features | cheap, pure | **warm-up / practice**, game/break phases |
89
90         **The scorer arbitrates.** It is the only place a decision is made, and it makes it by
*weighting columns*,
91         never branching. Today: `_gate_mask` (threshold) + `LogisticRegression` (learned
weights, ADR-004). The
92         confidence column on every signal lets the scorer discount an uncertain producer instead
of trusting a
93         binary verdict ? the small-N-safe way to fold in a deterministic-ish cue.
94
95         **Reconcilers audit.** After the scorer decides, the backward seam (§6) lints the
decision against the
96         persisted evidence and proposes corrections ? report-first, never silent.
97
98         **Four invariants make this safe to grow** (the load-bearing design claims):
99         1. **Signal, not verdict** ? every producer's output is a weighted feature; the decision
path has no
100         producer-specific branch. ? unbounded cues without combinatorial special-casing.
101         2. **Expensive detection ? cheap decision** ? frame-scope detection runs once and is
persisted; window/
102         session signals and the scorer are pure re-scoring (EXPERIMENT_HARNESS §2). ?
experiments are offline,
103         deterministic, zero production risk.
104         3. **Report-first reconciliation** ? corrections are recorded and reviewable, never
applied silently. ?
105         human-in-the-loop, auditable.

```

```

106     4. **Promote only on measured held-out lift** ? LOVO on the golden set + the
`run_regression.py` gate;
107     default OFF; rollback = config flip. ? no-regression evolution.
108
109     ---
110
111     ## 3. The forward seam ? injectable analyzers
112
113     Pluggable producers ? registered like `segmenter_registry`
(`backend/pipeline/segmenters/base.py:35`) ?
114     that run over the **per-frame-aligned signals already produced**: the WASB shuttle
trajectory
115     (`.frame/.visible/.x/.y` + velocity), the person boxes per frame
(`PersonDetector.detections_per_frame`,
116     `person_detector.py:256`), the net line (`rally_gate.estimate_play_axis`, camera-
agnostic), and the window
117     bounds. Analyzers are **pure** ? no I/O, no torch/cv2 ? mirroring `fusion_features.py`'s
discipline, so
118     their math is unit-testable in the GPU-free dev container.
119
120     ### 3.1 Contracts (JSON-serializable; proposed `backend/pipeline/analyzers/`)
121
122     ```python
123     @dataclass(frozen=True)
124     class Signal:                                # one scalar cue + how sure the producer is ?
never a verdict
125         value: float; confidence: float = 1.0    # persisted as `<producer>__<key>` (+
`_confidence`), [0,1]
126
127     @dataclass(frozen=True)
128     class AnalyzerEvent:                          # ? the `events` table on confirmation (§5)
129         type: str; timestamp: float; confidence: float = 1.0 # "net_hit"/"double_bounce"/?
; src-absolute s
130         player: Optional[str] = None; details: Dict[str, Any] = field(default_factory=dict)
131
132     @dataclass(frozen=True)
133     class BoundaryHint:                          # OPTIONAL deterministic proposal ? never auto-
applied
134         kind: Literal["start", "end"]; timestamp: float; confidence: float; reason: str
135
136     @dataclass(frozen=True)
137     class AnalyzerResult:
138         signals: Dict[str, Signal] = field(default_factory=dict)
139         events: List[AnalyzerEvent] = field(default_factory=list)
140         boundary_hints: List[BoundaryHint] = field(default_factory=list)
141
142     @dataclass
143     class WindowSignals:                          # read-only pre-produced bundle (no I/O inside
the analyzer)
144         start: float; end: float; fps: float; frame_width: float; frame_height: float
145         shuttle: List[TrajectoryPoint]           # in-window points (`.frame/.visible/.x/.y`)
146         persons: Dict[int, List[Tuple[int, BBox]]] # frame_idx ? [(track_id, box)]; {} when
person cue OFF
147         net: Tuple[str, float]                   # (axis, position): rally_gate auto-estimate or
config
148         frame_states: Dict[int, "ShuttleState"]   # frame_idx ? state; {} until the
shuttle-state detector (§3.5)
149         base_stats: Dict[str, float]              # motion stats already computed for the window
150
151     class WindowAnalyzer(ABC):
152         name: str; version: str                  # registry key + column prefix; version ?
provenance on change
153         feature_keys: List[str]; event_types: List[str] # declared columns/events (for
the harness)
154         @abstractmethod

```

```

155     def analyze(self, w: WindowSignals) -> AnalyzerResult: ...
156     ...
157
158     The three roles an `AnalyzerResult` can fill:
159
160     | Role | Sink | Consumed by |
161     |---|---|---|
162     | *(a) features* (with confidence) | `candidate_segments.stats` as `<analyzer>__<key>`
163     (+ `_confidence`) | the scoring model **and** the harness, as feature columns (§4) |
164     | *(b) events* | the `events` table (buffered on the candidate, flushed on
165     confirmation, §5) | history / highlights / audit |
166     | *(c) boundary hints* (optional) | stashed in `stats`; never auto-applied | the
167     reconciler (§6) / a future boundary-aware scorer |
168
169     ### 3.2 Worked example ? the SERVICE-FAULT analyzer (window-scope)
170
171     The owner's held-shuttle bug has a sibling: a **service fault** where the serve hits the
172     net and the
173     shuttle comes to rest in the net region. Pure trajectory, GPU-free, no person cue
174     needed:
175
176     - **Input:** the window's shuttle trajectory + the `rally_gate` net estimate.
177     - **Logic:** the shuttle **enters the net region** (`|coord ? net| < deadband`) **and**
178     its velocity
179     **decays to ? 0** there (comes to rest) ? high-probability net-hit service fault.
180     Confidence scales with
181     how cleanly the halt sits inside the net band.
182     - **Emits:** a `service_fault` **signal** (value = presence, confidence = halt
183     cleanliness) + a `net_hit`
184     **event** at the halt timestamp + a rally-**END** `BoundaryHint` at that timestamp.
185
186     Crucially: the analyzer does **not** drop the window. It *proposes* `service_fault`; the
187     scorer sees
188     `service_fault` + `service_fault_confidence` as two more columns and **learns**
189     whether/how much to
190     down-weight. The boundary hint is an audited proposal the reconciler may act on ? never
191     a branch.
192
193     ### 3.3 Injection point ? generalize `_enrich_candidate_stats`
194
195     The FusionSegmenter hook (`fusion_hybrid.py:70`) already turns "compute fixed features"
196     into per-window
197     enrichment. We generalize it to **run each enabled analyzer** ? additively, so the
198     existing
199     `net_crossings`/person paths are unchanged:
200
201     ```python
202     def _enrich_candidate_stats(self, cw, points, fps, frame_width, video_path, idx_cfg):
203         stats = super()._enrich_candidate_stats(...) # base motion keys
204         stats["net_crossings"] = ... # Gate-A signal (kept)
205         # ? person features when the person cue is on (kept) ?
206         w = self._window_signals(cw, points, fps, frame_width, video_path, idx_cfg)
207         for analyzer in analyzer_registry.enabled(idx_cfg): # NONE enabled ? byte-identical
208             to control
209                 res = analyzer.analyze(w)
210                 for key, sig in res.signals.items():
211                     stats[f"{analyzer.name}__{key}"] = sig.value
212                     stats[f"{analyzer.name}__{key}_confidence"] = sig.confidence
213                 stats.setdefault("_events", []).extend(e.__dict__ for e in res.events) #
214                 stats.setdefault("_boundary_hints", []).extend(h.__dict__ for h in
215                 res.boundary_hints)
216         return stats
217     ```
218
219
220
221
222
223

```

```

204     Flag-gated, **default OFF** via the existing `indexing.fusion.cue_flags` (a sibling
`analyzer_flags`
205     map keyed by `analyzer.name`). With every analyzer OFF, the loop is empty ? candidate
rows are
206     byte-identical to the substrate ? the no-regression invariant holds (§7).
`net_crossings` and the person
207     features are the **prototype analyzers** the seam retrofits first (proving it with zero
behavior change);
208     every new cue after that is a pure add.
209
210     ### 3.4 Session-scope analyzers ? the WARM-UP / PRACTICE example
211
212     Some structure is invisible at the window scope. **warm-up** looks exactly like a rally
in any single
213     window ? shuttle in flight, two players, net crossings, sustained motion ? which is
*why* it is a notorious
214     false-positive source. It is only distinguishable by **timeline structure**: a long
contiguous active
215     stretch *without* the rally?dead-time?rally cadence of scored play, usually (not always)
at the start.
216
217     So warm-up needs a producer that sees the **whole sequence of windows**, not one window.
That is a sibling
218     tier ? a `SessionAnalyzer` that runs once per video (a second pass, after per-window
enrichment) and emits
219     a labeled **span** plus a per-window **membership feature**:
220
221     ```python
222     @dataclass(frozen=True)
223     class Span:
224         kind: str                                # "warmup" | "practice" | "break" | "game" | ?
225         start: float; end: float; confidence: float; reason: str
226
227     @dataclass(frozen=True)
228     class SessionResult:
229         spans: List[Span] = field(default_factory=list)
230         # window_index ? signals merged back into THAT window's stats (e.g. {"in_warmup":
Signal(0.8, 0.6)})
231         window_features: Dict[int, Dict[str, Signal]] = field(default_factory=dict)
232
233     class SessionAnalyzer(ABC):
234         name: str; version: str
235         span_kinds: List[str]; feature_keys: List[str]
236         @abstractmethod
237         def analyze(self, windows: List[WindowSignals], session: "SessionContext") ->
SessionResult: ...
238     ```
239
240     **`WarmupAnalyzer`** (worked):
241     - **v0 (free, no new model): **find the *leading contiguous high-activity span** with
low inter-window
242     dead-time and no serve-reset cadence, over the persisted window timings we already
have. Emit
243     `Span(kind="warmup", ?)` + per-window `in_warmup` = overlap fraction of the window
with the span.
244     - **v1 (learned): **fed by the *shuttle-state cadence** (§3.5) / the **audio thwack**
cadence ? warm-up =
245     continuous hits *without* the point-ending pause rhythm; generalizes to mid-session
practice / between-
246     game knock-ups (drop the position prior).
247
248     **You framed the consumption exactly right: "so the scorer may consume it."**
`in_warmup` is a **feature
249     column**; the scorer *learns* to down-weight windows inside the span. **Warm-up is the
poster child for

```

```

250     no-special-casing**: the naive version is `if t < warmup_end: drop` ? a hand-coded
branch hanging off a
251     fuzzy boundary; a 20-second error **deletes a real opening rally**. As a weighted
feature, a wrong boundary
252     just slightly mis-weights a couple of windows the other cues rescue. The span is *also*
persisted as
253     evidence a reconciler (§6) may use to down-weight a whole cluster, but the primary path
is feature ? scorer.
254
255     **It is measurable today with no new labels:** the golden labels already treat warm-up
knock-ups / pre-serve
256     as **dead time = negatives** (Roora starts a rally at serve-contact; fusion §1 names
"warm-up knock-ups" as
257     a precision target). So an `in_warmup` feature is directly A/B-able on the 6 golden
videos for **precision**
258     lift ? that is the promotion gate.
259
260     ### 3.5 Consuming a richer detector ? SHUTTLE-STATE (frame-scope producer)
261
262     A natural new **detector**: a per-frame shuttle-state classifier ? `in_air` /
`racket_contact` /
263     `inactive`. Note the scope: it **processes frames ? a new per-frame signal**, like WASB,
so it is a
264     **frame-scope producer in the detection layer** ? its spec belongs in
265     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (it is a new `RallyCue`),
**not** a `WindowAnalyzer`.
266     This doc records **how the framework consumes it** once it exists:
267
268     - a **`WindowAnalyzer`** turns the persisted per-frame state into window features ?
`racket_contact_count`
269     (? shot cadence), `inactive_ratio` (? dead-time) ? that flow into the scorer as
columns;
270     - a **`SessionAnalyzer`** (warm-up, §3.4) reads its cadence;
271     - the **reconcilers** (§6) get *stronger* evidence: `trailing_dead_time` / `over_merged`
currently infer
272     inactivity from a velocity threshold; a learned `inactive` state is a more direct,
occlusion-robust
273     evidence source for "truncate / split here."
274
275     **Stage it (don't build the model first).** Most of the state is derivable for free from
signals we already
276     have ? `in_air` vs `inactive` ? WASB velocity + visibility; `racket_contact` ?
trajectory **direction
277     reversals** (fusion §4: "each reversal ? a shot") **and** the scaffolded **audio
thwack**
278     (`backend/pipeline/audio/hit_detector.py`, ADR-007's #1 cue). So **v0 = a derived
`shuttle_state` analyzer**
279     over existing signals (no new model); measure it. **The dedicated per-frame vision model
earns its keep
280     only where v0 hits a ceiling: WASB trajectory gaps** ? occlusion, motion blur, fast
smashes ? where there
281     is no `(x,y)` to threshold but an appearance model could still call `racket_contact`.
**Label blocker:** we
282     have **zero** per-frame state labels (golden labels are rally intervals); a dedicated
model needs weak/
283     distilled labels or a new (minors-gated) annotation task ? the same gate as feature-
level fusion. If you can
284     weak-label state from existing signals, prefer the derived analyzer until the gap-
filling case is proven.
285
286     ---
287
288     ## 4. How signals reach the scorer ? columns, not branches
289
290     The persisted `stats` columns are exactly what the harness already lifts.

```

```

`load_video_candidates(...,
291     feature_names=[...])` (`experiment.py:543`) reads each name out of the row's `stats`
JSON into a feature
292     column; `predict()` (`experiment.py:213`) gates/scores on those columns and merges. So a
new producer's
293     signal becomes a scorer input by **adding its column name to a `Variant`'s
`feature_names`** ? no new
294     scoring code, no branch:
295
296     ```python
297     # A learned Variant that lets the scorer weight the new signals (service-fault + warm-up
membership):
298     fusion_plus = Variant(name="fusion+analyzers",
299                           feature_names=[*SHUTTLE_FEATURES, "person__players_in_court",
300                                           "service_fault__service_fault",
"service_fault__service_fault_confidence",
301                                           "warmup__in_warmup",
"warmup__in_warmup_confidence"])
302     compare(golden_videos, CONTROL, fusion_plus)          # LOVO-honest B ? A on the 6 golden
videos
303     ```
304
305     The model **learns the weights** ? down-weight a window with a high-confidence
`service_fault` or high
306     `in_warmup`; up-weight ?2 in-court players + multiple net-crossings. **What the scorer
cannot see, it
307     cannot special-case; what it can see, it weights from data.** Same logistic regression
as ADR-004; we add
308     columns, never a net.
309
310     ---
311
312     ## 5. Persistence
313
314     - **Features ? `candidate_segments.stats`** (`database.py:240`). Namespaced
`<producer>__<key>`
315     (+ `__confidence`). Persisted **as a shadow** even when the served decision ignores
them, so any future
316     `Variant` re-scores offline (EXPERIMENT_HARNESS §5.2). `version`/`name` recorded for
provenance.
317     - **Spans ? window membership + a session record.** A `SessionAnalyzer`'s per-window
`in<span>` features
318     ride in each window's `stats` (so they re-score like any feature); the span list
itself (start/end/kind/
319     confidence) is a small per-video record (a JSON sidecar first; a table only if it
grows).
320     - **Events ? the `events` table** (`database.py:227`, `add_event(segment_id, ts, type,
player, details)`).
321     Its FK is to `play_segments`, which exist only **after** the AI handoff confirms a
candidate. So
322     analyzer events are **buffered on the candidate** (`stats["_events"]`) and **flushed**
to `events` in
323     Phase 4, when `__candidate_id` ? `segment_id` is linked (`trajectory_hybrid.py:283`).
Candidate-level
324     discoveries survive for offline re-scoring; the `events` table stays "events on
confirmed segments."
325     - **Corrections record** (§6) ? report-first means corrections are **recorded, never
silent.** Phase one:
326     a JSON sidecar shaped like `experiment_leaderboard.json` (one row per issue: segment,
kind, evidence,
327     before?after, reason, timestamp). When auto-apply is promoted, add a `corrections`
table via the
328     `PRAGMA user_version` migration ladder (`database.py:88`, `if version < 3:`) so
applied corrections are
329     queryable and reversible.

```

```

330
331 ---
332
333 ## 6. The backward seam ? the reconciliation pass (a linter for rally decisions)
334
335 Given **predicted segments + the persisted per-frame/per-window/per-session signals**,
detect where the
336 **DECISION contradicts the EVIDENCE** and propose corrections. This is the backward dual
of the analyzers:
337 analyzers add evidence *before* the decision; reconcilers audit the decision *against*
evidence *after* it.
338
339 ### 6.1 Contracts (proposed `backend/pipeline/reconcilers/`)
340
341 ```python
342 @dataclass(frozen=True)
343 class Issue:
344     segment_index: int; kind: str          # kind = rule id, e.g. "trailing_dead_time"
345     evidence: Dict[str, Any]              # the signal values that triggered it
(auditable)
346     before: Interval; after: Optional[List[Interval]] # None=drop . [iv]=corrected .
[iv1,iv2]=split
347     reason: str; confidence: float
348
349 @dataclass(frozen=True)
350 class Correction:
351     issues: List[Issue]; corrected: List[Interval]      # the FULL corrected set
(idempotent)
352
353 class ConsistencyRule(ABC):
354     name: str; precedence: int                # FIXED ordering; lower runs first
355     @abstractmethod
356     def check(self, seg: Interval, signals: "SegmentSignals") -> Optional[Issue]: ...
357
358     def reconcile(segments, signals, rules) -> Correction: ... # fixed precedence;
idempotent
359     ```
360
361 ### 6.2 Worked rules (fixed precedence)
362
363 | # | Rule | Evidence contradiction | Correction |
364 |---|---|---|---|
365 | 1 | `trailing_dead_time` | segment still "in rally" at its end, but the shuttle went
inactive earlier | **truncate** to the last-active frame *(the owner's misfire fix)* |
366 | 2 | `no_net_crossing` | `net_crossings < 1` ? nothing ever crossed the net | **drop**
(or flag) |
367 | 3 | `over_merged` | a dead-time gap **>*** `merge_gap` of shuttle-inactivity sits
**inside** one segment | **split** at the gap |
368 | 4 | `boundary_slack` | first sustained shuttle activity / first net-crossing is well
inside the start | **tighten** start to serve-contact |
369 | 5 | `inside_warmup` | the segment lies inside a high-confidence `warmup`/`practice`
span (§3.4) | **flag** (drop only under apply-flag) |
370
371 Each rule reads only **persisted evidence** (trajectory activity per frame,
`net_crossings`, person
372 occupancy, analyzer signals/events, session spans, shuttle-state) ? it is a
deterministic linter, not a
373 re-detector. As richer producers land (shuttle-state, warm-up), rules get **stronger
evidence**, not new
374 branches in the decision path.
375
376 ### 6.3 Report-first, apply-optional (non-negotiable)
377
378 `reconcile()` returns **both** an issues report **and** a corrected set. It **NEVER
silently rewrites**

```

```

379     (auditability + the repo's honesty/attribution working agreement). The pipeline either:
380
381     - **(default)** persists the report and **surfaces it for human review** ? production
output unchanged; or
382     - **(flagged, `reconcile.apply`)** applies the corrected set, **recording every change**
in the
383         corrections record (before?after, rule, evidence). Rollback = flip the flag.
384
385     **It is a `Variant` transformation, so its lift is A/B-measurable BEFORE auto-apply.** A
reconciler is a
386     pure post-`predict` stage: `reconcile(predict(...), signals, rules).corrected`. Wrapping
that as a
387     `RECONCILED` variant lets `compare(CONTROL, RECONCILED)` report the LOVO lift on the 6
golden videos
388     ([QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) corpus) ? promote to auto-apply **only**
on measured lift.
389
390     **Idempotent; fixed rule precedence.** `reconcile(reconcile(x)) == reconcile(x)`
(truncating an
391     already-truncated segment is a no-op); precedence is a fixed table so the result is
deterministic and
392     order-independent of input segment order.
393
394     ---
395
396     ## 7. No-regression discipline (identical to the harness contract)
397
398     The same four guarantees EXPERIMENT_HARNESS.md §6 makes ? applied to every producer and
the reconciler:
399
400     - **Default OFF.** Analyzers gate on `analyzer_flags` (default empty); the reconciler is
report-first
401     (apply gate default off). Merged code **cannot** change served behavior until
explicitly enabled.
402     - **Shadow-persist, then promote.** Signals/spans/issues are persisted even while
ignored, so a `Variant`
403     re-scores them offline; a signal becomes a *decision input* only when a Variant that
references it wins
404     LOVO. (EXPERIMENT_HARNESS §5.2's shadow mode.)
405     - **Promotion only via the eval gate.** `run_regression.py` against the golden
`baseline.json` (exit 1
406     blocks); the experiment leaderboard records the spec hash + per-video + LOVO deltas +
label-set hash.
407     - **Rollback = flip a config flag, never `git revert`.** The pinned `CONTROL` variant is
always present.
408
409     The keystone test (extends the harness's existing invariant): **all analyzer flags OFF +
reconciler
410     report-only ? candidate rows AND predictions byte-identical to control.**
411
412     ---
413
414     ## 8. How it composes with fusion P2
415
416     - **Window-analyzers plug into fusion P2's enrichment hook.**
`FusionSegmenter._enrich_candidate_stats`
417     (`fusion_hybrid.py:70`) is the exact injection point §3.3 generalizes; `net_crossings`
and the person
418     features become the first registered analyzers (no behavior change). New cues are add-
a-producer, per
419     ADR-007 / MULTI_SIGNAL_FUSION_PLAN §2.
420     - **Session-analyzers run as a second pass** over the per-window stats the same hook
persists.
421     - **The shuttle-state detector is a new fusion `RallyCue`** (frame-scope producer) ? it
belongs in the

```



```

422     fusion plan; this framework only *consumes* it (§3.5).
423     - **Reconcilers consume fusion P2's persisted signals** ? the same
`candidate_segments.stats` plus the
424     persisted trajectory; no new detection.
425     - Everything inherits fusion P2's control-reproducing defaults, so the fusion segmenter
with all flags off
426     still seeds and hands off identically to `wasb_hybrid`.
427
428     ---
429
430     ## 9. Trade-offs & risks
431
432     - **Deterministic vs learned signals.** A producer may emit a verdict-shaped signal
(`service_fault=1.0`),
433     but we persist it as a **feature**, so the scorer still arbitrates. Risk: a
confidently-wrong cue.
434     Mitigation: the **confidence column** + the **LOVO gate** ? a bad signal earns ? 0
weight, never promotes.
435     - **The temporal confound is the hard core.** Warm-up-vs-play, multi-court, and held-
shuttle all look
436     identical in any single frame / frequency band ? QUALITY_ITERATIONS' standing lesson:
"the dead-time
437     confusion is rally-vs-non-rally activity in the same place and frequency band... only
the learned model's
438     temporal features crack it." This is *why* warm-up is a **feature the model weights**,
not a heuristic
439     span-cutter, and why session-scope structure earns its own tier.
440     - **Scope discipline.** Put a producer at its *natural* scope: shuttle-state is frame-
scope (a detector),
441     not a window analyzer; warm-up is session-scope, not per-window. Mis-scoping forces a
producer to either
442     recompute detection or miss the structure it needs.
443     - **Ordering / idempotence (reconcilers).** Fixed precedence table + an idempotence test
(`reconcile?reconcile == reconcile`).
444     - **Signal, not verdict.** A producer that *decides* (returns keep/drop) re-introduces
special-casing;
445     reviews must reject "analyzer returns a boolean keep/drop."
446     - **Report-first auditability.** Reconcilers never silently mutate output; every applied
change is recorded.
447     - **Small-N feature discipline.** Few, scale/translation-invariant columns
(counts/ratios/presence, never
448     raw coords ? MULTI_SIGNAL_FUSION_PLAN §3.2); the model stays a logistic regression,
not a net.
449     - **Labels gate the learned producers.** Shuttle-state and feature-level fusion need
labels we don't yet
450     have; warm-up is measurable *now* because golden treats it as negatives. Know which
producers are gated.
451     - **Cost & licensing.** Pure analyzers are GPU-free; person- and shuttle-state-dependent
producers pull
452     torch / a second WSL pass ? gate behind their cue flags; all deps stay MIT/Apache/BSD
(ADR-002/005).
453     - **Event-table timing.** The buffer-then-flush wiring (§5) is the one non-obvious
integration detail.
454
455     ---
456
457
458     ## 10. Generality & research positioning
459
460     This is written as a **framework**, not a badminton patch. Read sport-agnostically it
is: *a registry-based,
461     multi-scope **signal-fusion** architecture for **weakly-supervised temporal
event/segment detection**, with
462     a learned arbiter and an **auditable post-hoc correction** layer.* The pieces that
generalize:
463

```

464 - ****Multi-scope producers ? one feature table ? a learned arbiter.**** The
 frame/window/session taxonomy and
 465 the "signal-not-verdict" rule are domain-independent. Any temporal problem with
****weak, multi-modal cues**
 466 at different time scales** (broadcast sports event spotting, surgical-phase
 recognition, call-center /
 467 meeting segmentation, activity logs) fits the same shape: add a producer, add a
 column, let the arbiter
 468 weight it, gate promotion on held-out lift.
 469 - ****Detection ? decision (persist once, re-score forever).**** Decouples the expensive,
 hard-to-reproduce
 470 perception from the cheap, deterministic decision ? the property that makes the whole
 thing
 471 experimentally honest and cheap to iterate.
 472 - ****Auditable reconciliation as a first-class layer.**** Most pipelines bury post-hoc
 fixes as silent
 473 if-branches; treating consistency-vs-evidence as a ***report-first, A/B-measured*** pass
 is the transferable
 474 methodological idea (and the human-in-the-loop / safety story).
 475 - ****The honest measuring stick.**** Interval-level, scarce ground truth + ****leave-one-
 VIDEO-out**** (group by
 476 match, never a held-out window from the same video) is the discipline any small-corpus
 temporal study
 477 needs to avoid flattering itself.
 478
 479 ****What is defensibly contributed vs. what is standard.**** Multi-modal ****late fusion****,
 weak supervision, and
 480 learning a linear/logistic arbiter over hand-designed features are all standard. The
combination offered
 481 here ? (a) a uniform producer-registry across ****three temporal scopes****, (b) a hard
****no-branch**** decision
 482 invariant that makes extension monotone, (c) a ****report-first, harness-measured
 reconciliation**** layer, and
 483 (d) the ****detection?decision persistence**** substrate that makes every experiment a zero-
 risk offline
 484 re-score ? is what would be written up, with rally detection as the motivating hard
 problem.
 485
 486 ****Honesty caveat (attribution working agreement).**** No external results or citations are
 asserted here. A
 487 real write-up still owes: a proper ****related-work survey**** (temporal action
 segmentation/detection, sports
 488 event spotting, multimodal fusion, weak supervision ? to confirm which claims are
 actually novel), ****strong**
 489 **baselines****, ****per-producer ablations****, and reporting on ****>1 sport / camera**** to
 substantiate the
 490 generality claim. Today the corpus is ****6 matches**** (QUALITY_ITERATIONS) ? enough for
 direction-checks and
 491 LOVO sign-tests, ****not**** enough to claim a general model; corpus growth is the highest-
 ROI unblock. Treat
 492 §10 as a ***positioning sketch***, not a results section.
 493
 494 > ****Update (2026-06-14): that related-work survey is now done ? see [§13](#13-prior-art-
 terminology--course-corrections-literature-review-2026-06-14).****
 495 > The honest finding: the ***skeleton*** (TAL/TAS + late fusion) is standard and must be
****cited, not claimed****;
 496 > the only **`novel-combination`** is the report-first reconciler-as-linter bundle; warm-up-
 as-a-soft-feature is
 497 > the thinnest genuine gap. §13 also lists the verified citations and the prioritized
 course corrections.
 498
 499 ---
 500
 501 **## 11. Phased execution (owner-gated; ONE PR per slice, off `master`, default OFF)**
 502

```

503 - **A0 ? this doc**, indexed in `docs/README.md`, cross-linked from the peers + ADR-007.
*(this
504     deliverable; docs-only.)*
505 - **A1 ? analyzer contracts + registry**
(`WindowAnalyzer`/`AnalyzerResult`/`Signal`/`AnalyzerEvent`/
506     `BoundaryHint` + `analyzer_registry`); generalize `_enrich_candidate_stats` to iterate
enabled analyzers
507     (none enabled ? byte-identical). Retrofit `net_crossings` + person features as the
first two analyzers.
508 - **A2 ? service-fault analyzer** behind `analyzer_flags.service_fault` (default OFF);
signal + `net_hit`
509     event + end boundary hint; shadow-persisted; offline `compare(CONTROL,
+service_fault)` LOVO.
510 - **A3 ? event flush wiring** (buffer analyzer events on candidates ? `add_event` on
confirmation).
511 - **A4 ? `SessionAnalyzer` tier + `WarmupAnalyzer` v0** (leading-span heuristic over
window timings),
512     default OFF; persists `warmup__in_warmup` + the span record; measured for
**precision** lift on golden.
513 - **A5 ? derived `shuttle_state` analyzer v0** (velocity + visibility + direction-
reversals (+ audio);
514     **no new model**); emits `racket_contact_count` / `inactive_ratio`; measured LOVO.
*(The dedicated
515     per-frame shuttle-state DETECTOR is a separate MULTI_SIGNAL_FUSION_PLAN `RallyCue`
item ? gated on
516     per-frame labels; build only if v0 is blocked by trajectory gaps.)*
517 - **R1 ? reconciler contracts + registry + `reconcile()`**
(`ConsistencyRule`/`Issue`/`Correction`) with
518     the worked rules; pure, idempotent, unit-tested; **not** wired to apply (report-only).
519 - **R2 ? reconciler as a `Variant` transform**: optional post-`predict` stage;
`compare(CONTROL,
520     RECONCILED)` LOVO; report-only in production.
521 - **R3 ? corrections persistence + opt-in apply**: JSON sidecar first, then the
`corrections` table
522     (`user_version < 3`); apply behind `reconcile.apply`; **rollback = flip the flag.**
523
524     Each PR is gated by the no-regression invariant ($7) + `run_regression.py`. Order is a
suggestion; A4/A5/R1
525     are independent and can be reprioritized on owner steer.
526
527     ---
528
529     ## 12. Verification
530
531 - **Offline, runnable in this dev container (no GPU/torch/cv2):** pure analyzer math on
synthetic
532     trajectories + person boxes (service-fault fires on a synthetic net-halt, silent on a
clean rally;
533     `WarmupAnalyzer` v0 finds a synthetic leading knock-up span; derived `shuttle_state`
matches a scripted
534     velocity/visibility sequence); the **no-regression invariant** (all flags OFF +
reconciler report-only ?
535     candidate rows + predictions byte-identical to control); the **reconciler
idempotence** test
536     (`reconcile?reconcile == reconcile`) and fixed-precedence determinism;
`compare(CONTROL, variant)`
537     deterministic on synthetic candidate rows.
538 - **Metrics ? adopt the field's (see [§13](#13-prior-art-terminology--course-
corrections-literature-review-2026-06-14)/C4):**
539     report **segmental F1@{10,25,50} + edit score** (TAS) and **tight/loose mAP@?**
(action spotting) alongside
540     temporal-IoU. tIoU alone is nearly blind to the over-segmentation the reconciler
exists to fix, so a
541     reconciler win ? or an IoU-chasing rule that shreds rally structure ? is invisible
under tIoU.

```

```

542 - **Machine-side (owner, GPU/WSL ? full pipeline can't run here):** persist analyzer
signals on a real
543 clip; `compare(CONTROL, +service_fault)`, `compare(CONTROL, +warmup)`,
`compare(CONTROL, RECONCILED)`
544 LOVO on the 6 golden videos; confirm the owner's **trailing-dead-time misfire**
truncates, the
545 **held-shuttle / net-hit service feed** is down-weighted, and **warm-up windows** lose
precision-mass;
546 spot-check `net_hit` events on confirmed segments. **Grade with §13/C4 metrics + per-
fold CIs (C1).**
547 - **Docs (A0):** links resolve; indexed in `docs/README.md`; cross-linked from
548 [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md),
[EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md),
549 [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md), and ADR-007.
550
551 ---
552
553 ## 13. Prior art, terminology & course corrections (literature review, 2026-06-14)
554
555 A multi-agent web-research pass surveyed the eight design patterns of this framework
against the
556 literature and **adversarially verified every citation** (real-paper check + closer-
prior-art hunt).
557 **Result: zero fabricated papers**; a few wrong author-lines / metadata slips were
caught and are
558 corrected below. This section is the honest "are we reinventing the wheel?" answer and
the actionable
559 course corrections ? all **owner-gated, NOT yet implemented** (they fold into the A*/R*
phases of §11).
560
561 ### 13.1 Verdict ? what is standard vs. what is ours
562
563 > **TL;DR.** The architecture's *skeleton* is standard and must be **cited, not
claimed**: candidate-windows
564 > + learned scorer = **Temporal Action Localization** (proposal-then-score); the
reconciler's
565 > over-segmentation fixes = **Temporal Action Segmentation** refinement/post-processing;
"signal not verdict /
566 > learned arbiter over cue columns" = **late / score-level fusion** realized as
**stacked generalization** ?
567 > ~20 years old and *already present in racket sports*. The **defensible delta is the
engineering/governance
568 > envelope**: detection?decision persisted as a re-score substrate + a **report-first,
A/B-lift-gated,
569 > idempotent, human-in-the-loop reconciler-as-linter** keyed on decision-vs-evidence
*contradiction* + the
570 > flag-gated no-regression promotion harness, on a deliberately tiny LOVO-by-match
golden set with an
571 > interpretable scorer. Of the patterns, only the **reconciler-as-linter** rated `novel-
combination`; the
572 > rest are `partial-overlap`. Park any novelty claim narrowly on (a) that reconciler
bundle and (b)
573 > **warm-up-as-a-soft-session-feature** (the one slice with genuinely thin prior art).
574
575 | Pattern (our framing) | Field's name | Closest verified prior art | Verdict |
576 |---|---|---|---|
577 | Candidate-windows + scorer + reconciler-as-refinement | Temporal Action Localization /
Segmentation + boundary refinement | ASRF (arXiv:2007.06866); Activity Grammars
(arXiv:2312.04266); BMN (arXiv:1907.09702) | partial-overlap |
578 | "Signal not verdict" ? learned arbiter over cue columns | Late / score-level fusion +
stacked generalization | Snoek 2005 (10.1145/1101149.1101236); Wolpert 1992; racket-sports A/V
rally-break fusion (CVIU 2009) | partial-overlap |
579 | Detection?decision, "persist once, re-score forever" | Frozen-backbone cached features
+ decoupled proposal/classification | S-CNN (arXiv:1601.02129); CBR (arXiv:1705.01180);
Decouple-SSAD (arXiv:1904.07442) | partial-overlap (? field default) |

```

580 | Report-first reconciler (contradiction ? audited correction, measured before apply) |
Model assertions + constraint-based repair ("ML linter") | Model Assertions (arXiv:2003.01668);
HoloClean (arXiv:1702.00820); TFDV (MLSys 2019); EventAnchor (arXiv:2101.04954) | ****novel-
combination**** |

581 | Scarce interval labels + tiny LR + LOVO-by-match | Point/timestamp-supervised TAL +
grouped (LOGO/LOSO) CV | SF-Net (arXiv:2003.06845); LOOCV-bias (arXiv:2406.01652); Snorkel
(arXiv:1711.10160) | partial-overlap |

582 | Multi-scope (frame/window/session); warm-up conditions per-window | Multi-
scale/hierarchical modeling + play-break + phase recognition | Tjondronegoro ICME 2003; GL-MSTCN
(10.1186/s12938-022-01048-w); MS-TCN (arXiv:1903.01945) | partial-overlap |

583

584 ****In-domain prior art that already hard-codes our cues**** (so our delta is the ***no-branch
discipline***, not the
585 cues): non-broadcast badminton rally start/end (See et al. 2024/2025), shuttle-state/hit
detection (MonoTrack
586 arXiv:2204.01899; Hsu/Yu/Cheng, Sensors 2024 24(13):4372), service-fault (Goh et al.,
Sensors 2023 23(24):9759),
587 tennis in-play state machines (10.1080/19346182.2013.819007), badminton point-
segmentation (Ghosh et al. 2018,
588 arXiv:1712.08714), ShuttleSet (arXiv:2306.04948).

589

590 **### 13.2 Terminology to adopt** (gloss our nouns with the field's on first use)

591

592 | Our term | Field's established term | Anchor |
593 |---|---|---|

594 | `RallyCue` / "producer" | ****labeling function**** / weak-supervision source; ****base
model**** in a stack | Snorkel; Wolpert 1992 |

595 | "Signal not verdict / no special-casing" | ****late / score-level fusion****; arbiter =
****level-1 combiner / meta-learner**** | Snoek 2005; Wolpert 1992 |

596 | ADR-004 LR scorer (`backend/eval/classifier.py`) | ****stacked meta-learner / level-1
combiner**** | Wolpert 1992 |

597 | "Detection?decision / persist once, re-score" | ****frozen-backbone cached features****;
****decoupled proposal/classification****; offline re-scoring | S-CNN; Decouple-SSAD |

598 | `WindowAnalyzer` (window-scope) | ****temporal action proposal**** features /
****actionness**** | S-CNN; BMN |

599 | Session-scope / warm-up | ****video-level temporal context / phase segmentation****;
****play-break**** outer layer; ****global-local**** | Tjondronegoro 2003; GL-MSTCN |

600 | `ConsistencyRule` reconciler | ****model/consistency assertions****; ****constraint-based
repair****; ****boundary refinement****; ****activity grammar / constrained decode**** | Kang 2020;
HoloClean; ASRF; Activity Grammars |

601 | "Report-first linter" | ****data validation / schema anomalies**** ("unit tests for data")
| TFDV; Deequ |

602 | "Over-merged / split" · "tighten boundary slack" | ****over-segmentation error**** ·
****boundary regression/refinement**** | MS-TCN; ASRF; CBR |

603 | "LOVO, group by match" | ****leave-one-group-out / -subject-out (LOGO/LOSO)**** grouped CV
| sklearn `GroupKFold`; Austin 2025 |

604 | Temporal-IoU metric | ****mAP@tIoU**** (TAL) + ****F1@{10,25,50} & edit score**** (TAS) +
****tight average-mAP**** (spotting) | Lea 2017; SoccerNet; Soares 2022 |

605

606 Keep the internal nouns; just gloss each on first use. Low-effort, high-leverage ?
reviewers map instantly to

607 30 years of literature and see we know the failure modes.

608

609 **### 13.3 What we can / cannot claim**

610

611 - ? ****The integrated governance envelope as a whole**** ? no found paper bundles all of:
report-first

612 contradiction detection over a ***persisted temporal decision*** + auditable structured
interval corrections

613 (truncate/split/drop/tighten) + ****per-rule A/B lift-gating before auto-apply**** +
idempotency/precedence +

614 human-in-the-loop + feeding a learned (value, confidence) scorer.

615 - ? ****"Signal not verdict" as a deliberate *design-philosophy delta***** vs. the in-domain
prior art that

616 hard-codes the same cues as if/else verdicts (state it as a philosophy delta, not a

capability they lack).

617 - ? ****Warm-up/practice as a soft session-scope feature**** ? thin prior art (only patents/coaching content surfaced); claim cautiously, validate from scratch.

618 surfaced); claim cautiously, validate from scratch.

619 - ? The ****fusion mechanism**** (multi-cue features ? light scorer for rally segmentation) ? ~20 yrs old in racket sports (CVIU 2009; Ghosh et al. 2018 learn an SVM over per-frame features ? smoothed rallies).

620 racket sports (CVIU 2009; Ghosh et al. 2018 learn an SVM over per-frame features ? smoothed rallies).

621 - ? The ****persist-once/re-score substrate**** as novel ? it is the ***default*** of the TAD field (datasets ship precomputed features so cheap heads re-train). Claim only the ***idempotent/versioned discipline***.

622 precomputed features so cheap heads re-train). Claim only the ***idempotent/versioned discipline***.

623 - ? The ****cross-check-and-correct insight**** ? Hsu et al. 2024 already cross-check shuttle-vs-swing streams; EventAnchor already does human-in-the-loop suggest-then-verify. Our delta is ***measured-before-apply + auditability***, not the cross-check.

624 EventAnchor already does human-in-the-loop suggest-then-verify. Our delta is ***measured-before-apply + auditability***, not the cross-check.

625 auditability*, not the cross-check.

626 - ? ****Static single-camera as a novel setting****, and ? the ****no-regression/flag-gated harness**** as a primitive (textbook MLOps) ? claim only "this discipline on a deliberately tiny golden set."

627 primitive (textbook MLOps) ? claim only "this discipline on a deliberately tiny golden set."

628

629 **### 13.4 Course corrections (owner-gated; prioritized)**

630

631 ****Tier 1 ? protect the honesty of "measured lift" on ~6 matches (do first):****

632 - ****C1 ? bias-aware LOVO.**** Standard leave-one-out on ~6 imbalanced groups induces a train-vs-heldout label anti-correlation that biases the estimate (a LR scored ~0.23 auROC on ***random*** data in the cited study).

633 anti-correlation that biases the estimate (a LR scored ~0.23 auROC on ***random*** data in the cited study).

634 Use class-balance-preserving folds (RL00CV) and ****report per-fold variance + CIs, never a single mean****, so the promotion gate can't promote noise. ***(Austin/Pe'er/Korem 2025, arXiv:2406.01652.)*** ? ****harness/eval.****

635 the promotion gate can't promote noise. ***(Austin/Pe'er/Korem 2025, arXiv:2406.01652.)*** ? ****harness/eval.****

636 - ****C2 ? calibrate each producer's confidence**** (Platt/isotonic) ***before*** the LR, so a CNN softmax, an audio onset, and an integer net-crossing count are comparable and the over-confident cue can't swamp the arbiter.

637 onset, and an integer net-crossing count are comparable and the over-confident cue can't swamp the arbiter.

638 ***(MoCaE arXiv:2309.14976; modality competition arXiv:2203.12221.)*** ? ****producers + scorer.****

639 - ****C3 ? out-of-fold (honest) stacking.**** When a producer is itself learned (shuttle-state), feed the LR its ***out-of-fold*** outputs; extend the group-by-match split to producer training. ***(Wolpert 1992.)*** ? ****detection?decision + harness.****

640 ***out-of-fold*** outputs; extend the group-by-match split to producer training. ***(Wolpert 1992.)*** ? ****detection?decision + harness.****

641 ****detection?decision + harness.****

642 - ****C4 ? report the field's metrics.**** Add ****segmental F1@{10,25,50} + edit score**** and a ****tight/loose mAP@***** sweep alongside temporal-IoU: tIoU is nearly blind to over-segmentation, so a reconciler win (or an IoU-chasing rule that shreds structure) is invisible under tIoU alone. ***(MS-TCN; ASRF; SoccerNet arXiv:1804.04527; Soares 2022 arXiv:2206.07846.)*** ? ****harness/eval + reconciler.****

643 mAP@*** sweep alongside temporal-IoU: tIoU is nearly blind to over-segmentation, so a reconciler win (or an IoU-chasing rule that shreds structure) is invisible under tIoU alone. ***(MS-TCN; ASRF; SoccerNet arXiv:1804.04527; Soares 2022 arXiv:2206.07846.)*** ? ****harness/eval + reconciler.****

644 IoU-chasing rule that shreds structure) is invisible under tIoU alone. ***(MS-TCN; ASRF; SoccerNet arXiv:1804.04527; Soares 2022 arXiv:2206.07846.)*** ? ****harness/eval + reconciler.****

645 arXiv:1804.04527; Soares 2022 arXiv:2206.07846.)* ? ****harness/eval + reconciler.****

646

647 ****Tier 2 ? architecture/discipline upgrades:****

648 - ****C5 ? resolve reconciler rules GLOBALLY**** (an activity-grammar / constraint-aware Viterbi decode), not as stacked local edits with fixed precedence (which can conflict / be non-idempotent across orderings). Model the structural priors (transition/duration) from data. ***(Activity Grammars arXiv:2312.04266; Constraint-Aware Decoding arXiv:2605.10149 ? read before citing.)*** ? ****reconciler.****

649 stacked local edits with fixed precedence (which can conflict / be non-idempotent across orderings). Model the structural priors (transition/duration) from data. ***(Activity Grammars arXiv:2312.04266; Constraint-Aware Decoding arXiv:2605.10149 ? read before citing.)*** ? ****reconciler.****

650 the structural priors (transition/duration) from data. ***(Activity Grammars arXiv:2312.04266; Constraint-Aware Decoding arXiv:2605.10149 ? read before citing.)*** ? ****reconciler.****

651 Decoding arXiv:2605.10149 ? read before citing.)* ? ****reconciler.****

652 - ****C6 ? brand the reconciler as "model/consistency assertions,"**** attach a per-correction confidence, and route only high-confidence corrections to auto-apply, the rest to a margin-ranked human queue; version the rule set as a TFDV-style owner-ratified anomaly schema. ***(Kang 2020; HoloClean ~90%**

653 route only high-confidence corrections to auto-apply, the rest to a margin-ranked human queue; version the rule set as a TFDV-style owner-ratified anomaly schema. ***(Kang 2020; HoloClean ~90%**

654 rule set as a TFDV-style owner-ratified anomaly schema. ***(Kang 2020; HoloClean ~90%**

```

repair precision ? ~10%
655     bad; Confident Learning arXiv:1911.00068.)* ? **reconciler + harness.**
656     - **C7 ? a few EARLY-fusion interaction columns** (e.g. `shuttle_in_air AND
person_two_sided`) ? a linear
657     arbiter can't represent AND/XOR; hand-crossed products recover it without leaving
"signal not verdict."
658     *(Snoek 2005.)* ? **producers + scorer.**
659     - **C8 ? represent a missing cue as an explicit `is_present`/uncertainty column, not a
silent 0** ? "cue
660     absent" ? "cue says no" for a linear arbiter. *(DBF arXiv:1511.03183.)* ? **producers
+ scorer.**
661     - **C9 ? keep warm-up/session-state a SOFT feature, never a hard pre-filter**, and add a
court/context
662     invariance probe: the static court co-occurs with rallies, so the scorer can learn
"court visible" instead
663     of "point in play" (action-context confusion). Evaluate on warm-up stretches that
share the court but lack
664     point cadence. *(arXiv:2103.16155; GL-MSTCN.)* ? **session-scope producers + scorer.**
665
666     **Tier 3 ? leverage external resources / higher-effort:**
667     - **C10 ? model cue CORRELATIONS** (Snorkel/MeTaL label-model baseline; strong L1/L2):
net-crossing,
668     two-sidedness, shuttle-in-air co-fire, so an unregularized LR double-counts the
bundle. *(arXiv:1711.10160.)*
669     - **C11 ? a temporal-smoothing pass** (Kalman/HMM/Viterbi over the per-window posterior,
or an MS-TCN
670     truncated-MSE anti-fragmentation penalty) before cutting intervals ? SmartTennisTV
hits F1 ~97.5% largely
671     via Kalman-smoothing noisy per-frame predictions. *(arXiv:1801.01430; MS-TCN.)* ?
**reconciler/scorer.**
672     - **C12 ? pull ShuttleSet/ShuttleSet22 + CoachAI** as an external benchmark + weak-
supervision corpus to
673     attack the ~6-match bottleneck. *("Less is More" arXiv:1903.00859; ShuttleSet
arXiv:2306.04948.)* ?
674     **harness/eval + producers.**
675     - **C13 ? benchmark the rule reconciler against a learned boundary refiner** (ASRF/CBR)
to *justify* staying
676     rule-based, and ground rules in MonoTrack-style physics/cadence as **soft features,
not hard rules** (hard
677     cadence rules would re-introduce special-casing and break on doubles/lets/warm-up).
*(MonoTrack
678     arXiv:2204.01899; CBR arXiv:1705.01180.)* ? **reconciler.**
679
680     ### 13.5 Pitfalls we are actively walking into (risk register)
681
682     - **Over-segmentation hidden by tIoU** ? the dominant error mode of cheap segment
decisions; a per-window
683     scorer can look good while fragmenting rallies. ? C4.
684     - **Late-fusion miscalibration / dominant-cue swamping** (over-confident shuttle CNN,
low-discrimination
685     audio). ? C2.
686     - **LOOCV distributional bias + tiny-N variance** ? promotion gate promotes noise
without CIs. ? C1.
687     - **Stacking leakage** (in-sample learned-producer outputs) ? C3; and **threshold-
tuning-on-test** ? keep
688     *every* tunable knob (windowing, merge/NMS, regularization, reconciler thresholds)
*inside* the LOVO loop.
689     - **Proposal-recall ceiling** ? the decision can't recover a rally the windowing missed;
measure window
690     recall (AR@AN) separately or you'll misattribute recall loss to the scorer.
691     - **Report-first decaying into silent rewrite** ? without per-rule lift gating + a
confidence gate +
692     idempotency tests, an over-eager rule systematically truncates correct long rallies.
693     - **Static-tripod transfer gap** ? broadcast cues (camera-switch, replays, crowd audio)
that drive the

```

694 tennis/soccer baselines *don't exist* on our footage; don't assume their F1s transfer.
695 - **Detector domain gap inherited downstream** ? WASB/TrackNet/MonoTrack are trained on
other rigs;
696 per-frame shuttle accuracy degrades off-distribution and every cue inherits that
error. **Quantify detector
697 accuracy on our footage before trusting any cue column** (currently unmeasured).
698
699 ### 13.6 Residual gaps & honesty notes
700
701 - **Fabrications: none** across ~40 citation checks. Corrected author-lines applied
above: CVIU 2009 =
702 **Liu, Huang, Jiang, Xing, Ye, Gao** (not "Kolekar"); SmartTennisTV = **Ghosh &
Jawahar**; Sensors 2024
703 shuttle-fusion = **Hsu, Yu, Cheng** *. Metadata fixes: L00CV-bias is **Science
Advances** (10.1126/sciadv.adx6976);
704 ShuttleSet (KDD) = arXiv:2306.04948 vs ShuttleSet22 = arXiv:2306.15664 (distinct); SF-
Net evaluates on
705 THUMOS14/GTEA/BEOID.
706 - **Read before citing** (verified to *exist*, not deep-read; both 2026, concurrent ?
don't void novelty but
707 are the nearest precedents): Constraint-Aware Decoding (arXiv:2605.10149, the
reconciler end-state analog)
708 and TennisExpert (arXiv:2603.13397, an audio-hits?cluster?filter rally pipeline). PCL-
Former
709 (arXiv:2507.06411) ? authorship unverified; cite without asserting authors.
710 - **Weakest-evidence claims (state plainly):** warm-up-as-soft-feature (novel by
absence of counterexample,
711 the high from-scratch-validation risk) and the measured-before-apply reconciler bundle
(also absence-based).
712 - **Not yet searched:** a recent badminton full-match frame-level rally/interruption
dataset (would relieve
713 scarcity); doubles-specific cue behavior; calibration for *count-valued* cues (net-
crossings); whether
714 RL00CV interacts badly with grouped folds at N=6 (the cited study is record-wise, not
group-wise).
715

12. user (2026-06-16T03:12:21.971Z)

1 # Owned-model implementation plan (roadmap) ? DRAFT for owner review
2
3 **Status:** PLAN / roadmap. No owned-model *implementation* starts until the owner
greenlights a specific
4 milestone. Paid Gemini stays the always-ready inference/fallback throughout (never a
training source).
5 **Grounded in:** `docs/archives/research/OWNED\_MODEL\_TRAINING\_STUDY.md`, the n=6 LOVO +
learned prototype
6 (`docs/archives/quality-iterations/2026-06-multivideo-lovo/`), and a verified strategy
workflow (2026-06-09).
7
8 ## Design principle (owner directive): *envision now, build later*
9 Every deferred piece (conversational-QA, on-device, multisport, the VLM) must be
**architecturally envisioned
10 now** so we never hit a mammoth refactor. We **design the seams** (data contracts,
export boundaries, the
11 event store) up front and **implement behind them** when each milestone is greenlit.
Punting a *feature* is fine;
12 punting its *seam* is not.
13
14 ## North Star (reframed per PR #61 review)
15 - **Immediate deliverable:** a **training framework + harness** that produces a model
with **very high precision/
16 recall for extracting highlights + rallies from a clip**. *Quality is the first-order
objective.*
17 - **Eventual product:** a **single, sport-agnostic, conversational** video model ?


```

upload a clip, then ask
18     contextual questions ("longest rally", "make a clip of all service faults", "rallies
over 20 shots").
19     - **On-device is LATER** (a model-compression goal so power users save upload
cost/hassle) ? definitely on the
20     roadmap, but **not** the immediate target. Don't over-index the strategy on it now.
21     - **The moat = the consented dataset + the eval harness, not the weights.**
22
23     ## Hard guardrails (non-negotiable ? "never corrupted")
24     - **Commercial-clean licensing for EVERY dependency ? code, data, AND model weights: MIT
/ Apache-2.0 / BSD only.**
25     (Ratifies the prior "Apache-2.0 only" to include MIT/BSD.) Reject viral/NC/custom-
restrictive: Gemma 3 (viral),
26     **Gemma 3n** (Gemma Terms follow derivatives ? even cleaner reject), Llama 4 (700M-MAU
cap), Commons-Clause repos.
27     - **Never train on paid-LLM (Gemini/OpenAI/Anthropic) outputs.** Also forbidden: public
grounding-CoT datasets that
28     were Gemini-generated (e.g. TVG-R1's 56K) ? we must originate our own CoT supervision
(a real, binding cost).
29     - **Exclude minors; per-user training data needs a separate explicit opt-in; split data
by match.**
30     - ? **Base-model pretraining provenance is unauditabile for *every* open VLM (incl.
Qwen3-VL).** Our "no paid-LLM
31     training" rule is enforceable at *our fine-tuning data* layer, not the base's
pretraining ? an **explicit owner
32     risk-acceptance** is required to use any open VLM. (Governed by the base's license,
not our pipeline.)
33
34     ## The model strategy (verified 2026-06-09)
35     **Framework ? base** (the Gemma-4-vs-ms-swift clarification): a *framework* is the forge
(runs SFT/GRPO, emits
36     weights); a *base* is the metal you forge. You pick one framework and feed it any clean
base.
37     - **Framework: ms-swift (Apache-2.0)** ? the one forge that does native **video + audio
+ timestamp-grounding +
38     GRPO/RFT with a custom video reward** in one stack. Keep **unsloth** (Apache-2.0 core;
*avoid its AGPL Studio UI*)
39     as the 8 GB-local SFT/compression tool for the deferred on-device goal. ? Point-in-
time snapshot ? re-verify at impl.
40     - **Base (when we reach the VLM): Qwen3-VL (Apache-2.0, uniform across sizes)** ? long-
video context, purpose-built
41     **timestamp emission** (second-level localization), multi-turn video chat. Re-confirm
the exact size's LICENSE
42     pre-release. **Gemma-4 stays AMBER and bench-only** ? its Apache grant is likely but
its Prohibited-Use posture is
43     unconfirmed (counsel needed), *and* its video/timestamp capability is disputed; not
the primary extractor.
44     **InternVL3 (MIT/Apache)** = clean fallback if a Qwen blocker appears. ? base swap is
cheap *only within
45     timestamp-capable Apache bases* (timestamp data-format + grounding-CoT are base-
specific).
46     - **START WITH THE TAL HEAD, NOT THE VLM** (the decisive call): fine-tuned VLMs **miss
strict temporal boundaries**
47     (best RL-post-trained video VLM ? R@0.7 50% on open-domain benchmarks; "coarse
regions, not accurate boundaries").
48     **Boundary accuracy IS the product.** The head (`rally_seq_proto.py`) is already de-
risked on our corpus (held-out
49     AUC 0.83?0.92), trains in minutes at ~4.5 GB on the 2070 Super ($0), is ~1M params
(trivially compressible later),
50     and **doubles as a pseudo-label teacher + the honest baseline the VLM must beat.** So:
**head now ? VLM later**,
51     strictly sequenced. When the VLM comes, prefer **two-stage** (VLM coarse-proposes,
head refines boundaries) over
52     wholesale replacement. *(Open: the VLM-ceiling number is open-domain; our amateur
footage may differ ? re-measure.)*

```

```

53
54     ## Conversational video-QA ? build the bookkeeping NOW, punt the model
55     **Punt the conversational *feature*** until extraction quality is solid (current LOVO
~0.583 ? "very high"; a chat UI
56     over a weak extractor exposes the weakness). **Reject end-to-end-VLM-answers-each-turn**
(expensive; bad at exact
57     counting; the strong teachers are license-forbidden). **But build the *seam* now** (the
brief requires it):
58     - **Architecture = "extract once, query many":** EXTRACTION (TAL head ? later VLM) emits
structured timestamped
59     spans ? a durable **per-video EVENT STORE** ? a **structured QUERY layer** (text-to-
SQL + ffmpeg clip-cut). The
60     cited queries ("longest rally", "all service faults", "rallies > 20 shots") are
**structured-numeric ? SQL, not a
61     VLM job** (SQL counts exactly where VLMs hallucinate).
62     - **Designed-now seam:** make the **per-video event-index schema a first-class data
contract** alongside the M0.1
63     corpus spine: per rally `{video_id(MD5), rally_ordinal, sport, start, end, shot_count,
ending_reason/fault_tags,
64     derived_clip_ref, provenance, confidence}`. **Reserve the event/fault-tag taxonomy
(e.g. `service_fault`) up front**
65     even if unpopulated. ? This is a real DB delta (a `PRAGMA user_version` migration):
`play_segments` today has
66     `video_id/start_time/end_time/duration/server/receiver/metadata` only; `highlights`
isn't a table yet ? so it's an
67     *extension with new columns + a highlights table*, not a no-op. **Punt** the NL
parser/agent until quality clears the bar.
68     - Paid Gemini may answer genuinely open-ended edge questions **at inference** over our
structured index (allowed ?
69     it reasons over our data, it is not trained on its outputs).
70
71     ## Multisport (single sport-agnostic model)
72     - **Model/framework: UNCHANGED.** One shared TAL head + one Qwen3-VL base, fine-tuned
across sports (sport = an input
73     feature / per-sport calibration). The collector + indexer are already sport-generic. ?
**ACTION: rename the repo off
74     "badminton" ASAP** (owner emphatic) ? consistent with this.
75     - **What multisport changes = DATA + DETECTORS (scale up, the accepted "good
problem"):** each new sport needs its own
76     labeled matches (split by match) + a **clean trajectory detector** to feed the head's
features. **The binding blocker
77     is per-sport detectors:** WASB (badminton) has the unresolved **C3** checkpoint-
provenance issue, and **clean MIT/
78     Apache/BSD shuttle/ball detectors for tennis/TT/pickleball/padel are not yet
identified.** Sequence by public-data
79     richness: **badminton ? tennis/TT ? pickleball/padel.**
80     - ? **Single-model-multisport-temporal-generalization is an OPEN HYPOTHESIS**, not
proven ? our corpus is badminton-only,
81     and the evidence once cited for it (RacketVision) was *refuted on verification* (it's
ball-tracking, not temporal). The
82     concept is sound on first principles (rally signal = a racquet-sport invariant) but
must be validated once a 2nd-sport
83     corpus exists. Don't cite RacketVision as evidence.
84
85     ## Language ? Python everywhere; on-device runtime is a deferred, layer-local export
shell
86     - **Training/eval = Python-locked, and that's fine** (PyTorch/ms-swift/ActionFormer);
offline/batch, no latency hot path.
87     - **Platform = no Python hot path to escape** ? the heavy bytes/GPU already run out-of-
process (ffmpeg, WASB-in-WSL, cloud
88     GPU per-job); the control plane is I/O-bound. **Reject a Go/Rust hot-path split** ?
single-digit-ms wins dwarfed by
89     ~40-min GPU jobs, and it would fracture the one codebase holding the QA-state + eval
harness + provider registry.
90     - **On-device (deferred) = the only place a 2nd language appears, and it doesn't dictate

```

```

our dev language:** train in
91     Python ? **export across the boundary** (ONNX/Core ML/ExecuTorch/GGUF) ? a thin device
shell (Swift/Kotlin/C++/Rust)
92     runs the artifact with zero Python. **The one concrete constraint now: train with
export-clean ops** so the compression
93     target stays reachable without re-architecting. (On-device base stays on the
Qwen3-VL-4B Apache lineage ? *not* Gemma-3-270M.)
94     - Escalation order if a real Python CPU hot path ever appears: free-threading ?
numpy/vectorize ? Cython/PyO3 ? only then a
95     separate service. **None exists today.**
96
97     ---
98
99     ## Phase 0 ? Foundation (start now, per-milestone greenlight; $0, local, no cloud/DPA)
100
101     ### M0.1 ? The labeled-corpus spine **+ the event-index contract** (the moat) .
*greenlight item*
102     Canonical JSONL row per rally (corpus) **and** the queryable per-video event-index
schema (the conversational seam)
103     ? same row, wired to a store. Fields: `{video_id(MD5), rally_ordinal, sport, start, end,
ending_reason/fault_tags,
104     shot_count, label_type, source(human|pseudo|vendor), consent/training_opt_in, split(BY
MATCH), trajectory_ref,
105     labeled_window, derived_clip_ref, confidence}`. 3 transcoders (heuristic/LogReg . TAL
ActivityNet-JSON . VLM clip?QA),
106     one scorer spine (`metrics.match_intervals`). Authored in `sports-data-collector`
(system-of-record; extends PR #13).
107     **Reserve fault/event taxonomy now.** Acceptance: round-trip test; split-by-match
enforced; **no Gemini-sourced rows in any training view**.
108
109     ### M0.2 ? Grow the corpus . **DONE (n=6), running** . *owner action + my scoring loop*
110     6 distinct matches labeled (3 singles incl. Mahadevpura, 1 doubles, 2 multi-court); per-
match + LOVO recorded; the
111     "contrast" data-quality metric validated. Owner adding 1 more varied multi-court-club
match ? margin. **Bottleneck =
112     data + calibration, not method.** Keep growing per new sport.
113
114     ### M0.3 ? Compliance cleanup (**hard blocker ? do before any training run**) .
*greenlight item*
115     - **(i) PURGE the Gemini-silver TRAINING path** ? the live no-paid-LLM-training
violation is
116     **`backend/eval/distill_local.py`** (it trains a classifier on Gemini-silver CSV
labels ? `output/local_rally_model.json`)
117     and the threshold distillation in `calibrate_local.py`. (NB:
`backend/eval/training.py::label_candidates` is a *pure,
118     source-agnostic* function ? there is no store by that name; the earlier
"training.label_candidates" pointer was wrong.)
119     Nothing Gemini-trained is in git (artifacts live under gitignored `output/`), but the
path that regenerates them is one
120     command away ? retarget to human + open-teacher (own-model/WASB) labels; Gemini =
eval/reference/fallback only.
121     - **(ii) WASB C3 ? first-class action item (owner):** the academic-data checkpoint
provenance is a **commercial-ship
122     blocker** that a clean head does NOT launder. Resolve via **own-trained weights** or a
clean MIT/Apache detector swap.
123     ? **own-weights is an un-scoped NEW workstream, not ROADMAP Phase 4:** per ADR-002 our
temporal labels can't train the
124     shuttle detector (it needs per-frame coordinate labels we don't have). **Multiplies
per sport** (incl. any TT-on-WASB use).
125     Carry as a tracked blocker that **gates SHIP** (not Gen-0 research).
126
127     ### M0.4 ? Gen-0 TAL harness (**FIRST-CLASS ACTION ITEM** per owner) . *greenlight item*
128     Productionize `rally_seq_proto.py` ? a **one-command train?eval?ship-gate** harness over
**ActionFormer (MIT, 1:1
129     rally=instance)** and/or **ASFormer (MIT, TAS)** (use `sj-li/MS-TCN2` if MS-TCN++, *not*

```

```

the Commons-Clause repo),
130     reusing `metrics.match_intervals`. **Define the ship-gate target up front** (open item):
a FLOOR of "beat the honest n=6
131     heuristic LOVO **0.480**" is necessary but not sufficient ? set an explicit **F1@tIoU**
target for "very high P/R" (e.g.
132     F1@tIoU=0.5 for highlights, tighter for precise rally cuts) before declaring victory.
Trains on the GPU/WSL box.
133
134     ---
135
136     ## Phase 1 ? Gen-0 ship decision (gated on n?5 + M0.4)
137     Run the A/B; if the learned Gen-0 clears the ship-gate with per-video calibration, it
becomes the rally/highlight
138     extractor (local, MIT). If not, the gap (more data/features) is the next signal ? no
overfitting to declare victory.
139
140     ## Phase 2 ? Gen-1 / Gen-2 (gated on platform scaffolding P0?P4 + healthy corpus +
explicit go)
141     - **Gen-1:** concatenate fully-local cue channels onto the head. ? **Audio is NOT a
promising cue here** ? our E1 probe
142     (PR #35) found it weak (~2.2x discrimination, AUC ~0.6) and the foreground-audio idea
also failed this session; keep it
143     parked, not "in reserve." Pose/court is the more credible #2 cue (license/cost-vet
first).
144     - **Gen-2 (VLM):** ms-swift fine-tune of **Qwen3-VL** (SFT cold-start ? GRPO/RFT with a
temporal-IoU-vs-golden reward,
145     our own CoT supervision). **Two-stage** (VLM proposes, head refines). Cloud train on
Modal (no-access) / RunPod-Secure.
146     - **Hard pre-serving gates (unmeasured, load-bearing):** the **$/call crossover** (self-
hosted VLM vs Gemini) **a real
147     VRAM/token + cost benchmark** on rented A100/L40S (ms-swift multimodal-GRPO examples
use 6?8 GPUs ? **budget is the
148     largest open risk and is currently unquantified**). Don't stand up an owned serving
stack until these clear.
149
150     ---
151
152     ## Decisions / status
153     | # | item | status |
154     |---|---|---|
155     | 1 | Greenlight Phase-0 milestone(s) ? rec: **M0.1 corpus+event-index**, then **M0.4
harness** | awaiting owner |
156     | 2 | WASB C3 ? first-class action item (own-weights / clean swap) | **owner: yes** ? |
157     | 3 | Ratify licensing **MIT / Apache-2.0 / BSD** (code+data+weights) | **owner: yes** ?
|
158     | 4 | Collector = corpus + event-store system-of-record | **owner: confirm** ? |
159     | 5 | Rename repo off "badminton" | **owner action** ? |
160     | 6 | Conversational-QA: build seam now, punt model | **rec accepted** ? |
161
162     ## Open risks / unknowns (none block Phase-0 research)
163     Budget/cost to reach Gen-2 **unquantified** (largest risk) . ship-gate "very high P/R"
target undefined . per-sport clean
164     trajectory detectors not yet identified . base-model pretraining provenance unauditale
(risk-acceptance) . real VLM
165     VRAM/token cost . single-model-multisport-temporal generalization unproven . DPA + no-
train (minors-excluded) for cloud
166     training . on-device 4B latency on 8 GB unmeasured.
167
168     ---
169     *Design-only and reversible. Nothing starts until you greenlight a milestone. Sequencing
discipline holds: the committed
170     next platform slice remains the async job model (single-tenant); only the M0.1?M0.4
foundation is safe to start.*
171

```

13. user (2026-06-16T03:12:22.222Z)

```
1      # Rally-detection quality iterations ? living log
2
3      **Purpose.** A running, reverse-chronological log of how we evolve rally-detection
quality: each
4      iteration's hypothesis, what shipped, the measured (or expected) effect, and the lesson.
This is the
5      **live** continuation of the historical trail in
6      [`archives/quality-iterations/`](archives/quality-iterations/README.md) ? per the
archive policy
7      (`CLAUDE.md`), **terminal-state** (DONE/REJECTED) work gets a dated subdir there; this
file stays at
8      `docs/` top level and keeps growing, then its entries graduate to the archive when they
conclude.
9
10     **Methodology (non-negotiable).** Every quality change rides the experiment harness
11     ([EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)): new cues land **flag-gated, default
OFF**; they're
12     graded **offline against golden labels** (per-video + LOVO-honest); promotion to default
happens **only
13     on measured lift** through the `run_regression.py` gate; **rollback = flip a config
flag, never
14     `git revert`. **Cues are designed in
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md); scoring
15     in [EVALUATION.md](EVALUATION.md); the 2-layer model in
16     [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md).
17
18     **Golden corpus = the 6 labelled matches** in `output/human_lovo_manifest.json` (owner-
confirmed
19     2026-06-13: use all 6 ? `adarsh_avi_singles, kushagra_singles, largetest_doubles,
gbaaddy,
20     `testlarge_short, mahadevpura_singles`). The model trained on these is what we run
alongside the existing
21     detector on NEW footage for the side-by-side review.
22
23     ---
24
25     ## Baselines that drive strategy (IoU 0.5 vs human golden labels)
26
27     | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |
28     |---|---|---|
29     | Heuristic (velocity windowing) | 0.657 | 0.157 |
30     | Gen-0 owned head (LOVO, n=6) | 0.744 | 0.561 |
31     | Two-stage WASB+Gemini refine | 0.83 | ? |
32     | Gemini wrapper (`gemini-flash-latest`) | 0.93 | 0.66 |
33
34     **Reading:** clean footage is commoditized (serve via Gemini); **multi-court drops the
commodity wrapper
35     to 0.66 ? that is the owned model's ship target. **Baselines tracked in
`eval_baselines/`. The owner's
36     observed false positive ? a player **holding/flicking** the shuttle between points
counted as a rally ?
37     is the precision bug the fusion gates target.
38
39     ---
40
41     ## Current track (live) ? newest first
42
43     ### Served-side Gate-A measurement on the LIVE FusionSegmenter ? the independent run
behind the #151 revert *(2026-06-14)*
44     The owner-gated follow-up flagged by the (now **REVERTED**) SERVED-DEFAULT entry below.
An **independent** served-side check ?
45     separate from, and corroborating, the [#148](https://github.com/avidullu/badminton-
highlight-indexer/pull/148) litmus + #151's
```

```

46     run ? of whether the Gate-A default flip (PR #146) regresses the live path. **Verdict:
it does ? the flip was reverted in
47     [#151](https://github.com/avidullu/badminton-highlight-indexer/pull/151); this entry is
the detailed measurement.** Driven the
48     **real** served seams (`_enrich_candidate_stats` ? `net_crossings`,
`_select_for_handoff` ? the gate) over each golden video's
49     **cached** WASB trajectory, at the **production windowing** `process_video` actually
uses
50     (`velocity_thresh=0.02, merge_gap=2.0, min_window=2.0`) ? no WASB re-run (the runner has
no cache-hit early-return), no GPU, no
51     Gemini. Config loaded exactly as production (`load_config()`); segmenters from the
registry. Repro:
52     `scratch/served_gateA_verify.py` + `scratch/served_gateA_sweep.py` (artifacts in
`output/gateA_served_verify/`); the standing
53     productionized guard is `backend/eval/served_gate_a.py` (#148).
54
55     - **(a) No-regression invariant CONFIRMED on real served code.** At `min_crossings = 0`,
`fusion_hybrid`'s seeding **and**
56     handoff set are **byte-identical to `wasb_hybrid`** on all 6 golden videos (same
candidate windows; handoff = all windows =
57     the base seam). The "fusion adds nothing when the gates are off" property is now
verified on the *live* served path, not
58     only by code review. ?
59     - **(b) The +0.074 win does NOT transfer to the served windowing.** The offline
promotion (+0.074, 6/6, p=0.031) was measured
60     on the eval harness's **recall-oriented *proposal* windowing** `(0.005, 1.0, 0.5)`,
where CONTROL **over-segments**
61     (segment-count ratio 1.68) and Gate-A's job is to clean up that fragmentation (?
1.01). The **served path uses a coarse
62     windowing** `(0.02, 2.0, 2.0)` that already merges aggressively (low over-
segmentation), so Gate-A has little FP to remove
63     and instead trims **real** rallies whose WASB trajectory registered only 0?1 net
crossings. **Candidate-handoff-level F1
64     (pre-Gemini, IoU ? 0.5) on the served windowing, n=6:**
65
66     | min_crossings | mean F1 | precision | recall | real-rally windows lost vs off | ?F1
vs off |
67     |---|---|---|---|---|---|
68     | 0 (Gate-A OFF = `wasb_hybrid`) | 0.300 | 0.317 | 0.290 | 0 | ? |
69     | 1 | **0.303** | 0.327 | 0.288 | 1 | **+0.003** |
70     | **2 (shipped default, PR #146)** | 0.292 | 0.337 | 0.261 | **12** | **?0.008** |
71     | 3 | 0.269 | 0.338 | 0.226 | 28 | ?0.031 |
72
73     At the shipped `mc=2`, mean F1 goes **0.300 ? 0.292 (?0.008)** and Gate-A drops **12
windows that match golden rallies**
74     (net_crossings almost all exactly 1.0), concentrated in doubles (`targetest_doubles`
?0.051). All deltas are **tiny / within
75     n=6 noise**, so this is not a catastrophic regression ? but `mc=2` is **not** a
served-side *improvement*, and the
76     held-shuttle "guard" is removing genuine rallies, not just feeds.
77     - **Served relevance (honest).** A dropped candidate window **never reaches Gemini**, so
each lost real rally is an
78     **unrecoverable served recall loss**; the precision "gain", by contrast, is partly
what Gemini already does at the handoff ?
79     so the served-side tradeoff is at best neutral and plausibly net-negative, i.e. the
candidate-level ?0.008 likely
80     *understates* the served downside.
81     - **Root cause = eval ? serve windowing.** The harness promotes/refuses variants on a
windowing the production path does not
82     serve. Gate-A's value is conditional on over-segmentation; the coarse served windowing
doesn't over-segment, so the gate
83     inverts. [[decision-rigor-norm]]: the foundational issue (harness windowing ? served
windowing) outweighs the single
84     proposal-windowing number.
85

```

```

86      **Outcome ? REVERTED (see the entry below).**
[151](https://github.com/avidullu/badminton-highlight-indexer/pull/151)
87      restored the pre-#146 state: `config.json` `indexing.default_segmenter` ? `gemini` (the
`indexing.fusion` block removed) and
88      `FusionConfig.min_crossings` default **2 ? 0** ? Gate A is now **OFF by default but
still opt-in**
89      (`indexing.fusion.min_crossings: 2`), the persisted-superset invariant intact. Our
independent numbers match #151's (mean
90      served ?F1 ?0.008, ~12 rallies dropped, 2/6 losses). If Gate A is ever opted back in on
the served path, `mc=1` is the
91      least-bad operating point (?flat F1, 1 rally lost). **Open follow-up:** **align the
experiment-harness windowing with the
92      served windowing** (or re-run promotions on the production windowing) so a harness lift
actually predicts served behaviour ?
93      the real lesson here. The offline ablation + unit tests remain valid for what they
measured.
94
95      ### PERSONALIZATION Phase-0 quality checkpoint ? the eval/safety RAILS for the per-user
feature *(2026-06-14)*
96      The owner adopted the **HYBRID** posture: personalization as a **flag-gated FEATURE on a
generalization-first baseline**,
97      paired with a **contribution flywheel** so per-user gains and the shared baseline
compound together (full design in
98      [PERSONALIZATION_PLAN.md](PERSONALIZATION_PLAN.md), merged
[144](https://github.com/avidullu/badminton-highlight-indexer/pull/144)).
99      This checkpoint landed the three **honest-eval RAILS** that keep a per-user/small-N
tuning layer from promoting noise or
100      silently regressing a user:
101
102      - **Per-user promotion gate ? [141](https://github.com/avidullu/badminton-highlight-
indexer/pull/141)**
103      (`backend/eval/per_user_gate.py`): a **min-N floor** + **structural-guard exemption**
? per-user/small-N tuning can't
104      promote noise and can't disable a structural guard (the proven Gate-A held-shuttle
guard stays on regardless of a thin
105      user corpus).
106      - **Held-out-USER learning curve ? [142](https://github.com/avidullu/badminton-
highlight-indexer/pull/142)**
107      (`backend/eval/per_user_eval.py`): turns the vague "<K videos to superior quality"
claim into a **falsifiable release
108      number** (the `smallest_k_with_lift` measured on held-out users).
109      - **Per-user no-regression guard ? [143](https://github.com/avidullu/badminton-
highlight-indexer/pull/143)**: a
110      retrain/baseline-upgrade can **never make a user worse than their incumbent** (frozen
per-user CONTROL; a regressing
111      candidate is refused).
112
113      **Read HONESTLY:** these are the **eval/safety rails** for the personalization feature ?
**not** new measured
114      rally-detection wins. No F1 moved this checkpoint; what landed is the discipline that
lets per-user tuning ship without
115      regressing the baseline or over-claiming on n?6-per-user data.
116
117      **Locked owner decisions (2026-06-14):** identity = *design for the north-star (multi-
user/cloud), implement for current
118      use (single-user local)*; the inline per-user classifier sits **behind a cloud-load
interface** (local store = one impl);
119      **`min_n` is a PROMOTION-confidence floor, NOT a service/segmentation gate** (a one-
video user is ALWAYS fully segmented by
120      the baseline batteries); **free-tier videos ARE pooled** into the shared baseline
(consent baked into free terms,
121      minors-exclusion enforced at the pooling query); and the previously-pending structural-
guard sub-decision is now LOCKED ?
122      **Gate-A-only structural + `min_n = 5`**.
123

```

```

124     ### Fusion P3 + P4 MEASURED on the n=6 golden corpus ? person = no lift, audio =
promising (not promotable) *(2026-06-14)*
125     The GPU person-feature extraction finished all 6 golden videos (`fusion_golden.py` ?
`output/*_golden_features.json`, the
126     replayable substrate); the CPU LOVO drivers then measured the lift with **honest stats
(C1 ? bootstrap 95% CI + paired sign
127     test; never a bare mean at n=6)**:
128
129     | Variant (honest LOVO, IoU ? 0.5) | F1 | ?F1 vs prior | 95% CI | sign test | CI clears
0 |
130     |---|---|---|---|---|---|
131     | shuttle-only | 0.307 | ? | ? | ? | ? |
132     | + person (P3) | 0.286 | **?0.021** (?6.7%) | [?0.165, +0.162] | 2W/4L (p=0.688) |
**No** |
133     | + person + audio (P4 incr.) | 0.366 | **+0.079** (+27.7%) | [?0.054, +0.203] | 4W/2L
(p=0.688) | **No** |
134
135     **Read honestly:** neither cue clears the promotion bar (both CIs span 0, both sign
tests n.s.) at n=6 ? **both stay
136     default-OFF** (promote-only-on-lift). **Person** adds noise (no court polygons ?
`players_in_court` counts spectators);
137     per-video it's a wash (testlarge +0.40, gbaaddy ?0.22). **Audio** is the more promising
lead ? +0.079 / +27.7%, wins 4/6
138     (largetest +0.27, mahadevpura +0.27), and shuttle+person+audio (0.366) lands **above**
shuttle-only (0.307) ? but n=6 can't
139     confirm it. **Next:** more matches (raise N until the sign test can resolve ~+0.08), and
measure audio on shuttle-only
140     *directly* (without the person handicap). Recorded in
`eval_baselines/experiment_leaderboard.json`. **Negative/uncertain
141     results are kept on purpose ? this is the honest ablation backbone for the paper aim.**
142
143     ### Gate-A served-default flip ? **REVERTED (eval?serve regression)** *(2026-06-14)*
144     **The #146 flip of Gate-A (net-crossings ? 2) to the live served default was REVERTED.**
A served-side check ?
145     the repeatable litmus in [#148](https://github.com/avidullu/badminton-highlight-
indexer/pull/148)
146     (`backend/eval/served_gate_a.py` + `tests/test_served_gate_a_regression.py`),
corroborated by an independent
147     served run ? found Gate A is **not a win on the served path**: on the production
(COARSE) windowing it is
148     **neutral-to-negative, mean served ?F1 ?0.008**, dropping ~12 real rallies across the
golden 6 (2 losses, 0
149     clear wins). It is a **confirmed served regression**, so the owner directed a revert to
restore the known-good
150     state.
151
152     **Why eval ? serve (the honest root cause).** The +0.074 below is a property of the
**offline harness's
153     permissive PROPOSAL windowing** (`velocity_threshold=0.005, merge_gap=1.0,
min_window=0.5`), which emits many
154     short candidate windows ? leaving abundant held-shuttle junk for Gate A to cut. The
**served COARSE windowing**
155     (`velocity_threshold=0.02, merge_gap=2.0, min_window=2.0`) already merges flight bursts
into a few long windows
156     and drops short ones (served over-seg ratio already **< 1**, i.e. under-segmentation),
so there is little junk
157     left to cut and the gate instead occasionally trims a borderline true window ? slightly
negative ?F1. The gate
158     *itself* is correct (#148's control test reproduces +0.074 / 6-of-6 / p=0.031 exactly
under proposal
159     windowing); the **operating point is what differs**. This is the canonical eval?serve
trap.
160
161     **What the revert restores** (back to the pre-#146 no-regression state):
162     - `config.json` ? `indexing.default_segementer` **`fusion_hybrid` ? `gemini`** (the prior

```


served detector); the

163 explicit `indexing.fusion` block added by #146 is removed.

164 - `backend/config/models.py` ? `FusionConfig.min\_crossings` default ****2 ? 0**** (Gate A ****OFF**** by default).

165 - Tests (`tests/test\_config.py`, `tests/test\_fusion\_hybrid.py`) restored to the Gate-A-OFF defaults; the

166 ****persisted-superset invariant is preserved**** (gating only ever filtered the AI handoff, never persistence).

167 - `eval\_baselines/experiment\_leaderboard.json` ? the #146 `promotion` record carries a `reverted (eval?serve)`

168 note.

169

170 ****Gate A is NOT deleted ? it remains opt-in**** (`indexing.fusion.min\_crossings: 2`) for footage with more

171 held-shuttle/service-feed false positives than the golden 6, and the offline ablation evidence below stands

172 (it is a real proposal-windowing win). The path back to serving it is to ****align the served windowing toward**

173 the proposal operating point** and re-run the #148 litmus to confirm the lift transfers. The #148 litmus is

174 now the standing guard: it pins the honest measured behaviour (served ~neutral; +0.074 only under proposal

175 windowing) and trips if a future windowing change silently revives or breaks either fact.

176

177 **### Acting on the first A/B: Gate-A PROMOTED + learned variant FIXED + GX010125**

corroboration *(2026-06-14)*

178 ****Gate-A (net-crossings ? 2) ? PROMOTED (first cue promoted via the harness).**** The measured win

179 (?F1 +0.074, ****6/6 golden matches, sign-test p=0.031****, CI [+0.042,+0.104]; over-seg 1.68?1.01 ? the

180 [first-A/B archive](archives/quality-iterations/2026-06-first-ab-experiment/README.md)) clears the honest

181 bar. ****Recommended config:**** `indexing.fusion.min\_crossings = 2` on the fusion path (the eval-only Gate A

182 folded into production, MULTI_SIGNAL_FUSION_PLAN §3.1). ? The offline win is on the recall-oriented *proposal*

183 candidate set; flipping the ****live served default**** stayed owner-gated at the time of this entry. ****Update:** the

184 flip was attempted (#146) and then REVERTED** ? on the served COARSE windowing Gate A is neutral/negative

185 (?0.008), an eval?serve regression; see the REVERTED entry above. Gate A stays opt-in

186 (`indexing.fusion.min\_crossings: 2`), default OFF. ****Corroborated on real footage:**** offline

187 re-score of GX010125_1080p's *cached* trajectory (no GPU, isolated `output/GX010125\_run/`) ? 45 candidates ?

188 CONTROL 45 rallies vs ****Gate-A 27****; Gate-A dropped ****18 low-crossing windows (~1.7 min)**** of likely

189 held-shuttle/feed FPS (no golden labels for this footage ? divergence, not F1).

190

191 ****Learned variant FIXED ? [#130](https://github.com/avidullu/badminton-highlight-indexer/pull/130).**** The

192 first A/B's learned logistic lost to CONTROL and zeroed `testlarge` ? a fixed-0.5-cutoff artifact. Added

193 additive/default-OFF `Variant.tune\_threshold` (operating point chosen on the ****train**** fold) + `calibrate`

194 (Platt/isotonic, C2). Validated on the n=6 golden set: ****mean F1 0.307 ? 0.423 (+0.116)****, testlarge

195 0.000 ? 0.300, and the learned variant now ****beats CONTROL (0.388)****. Threshold-in-LOVO is the big lever;

196 calibration is available too.

197

198 ****Next:**** a real fusion-path run to confirm Gate-A served-side; per-video output layout

199 (PER_VIDEO_OUTPUT_LAYOUT.md) so each video's artifacts are self-contained.

```

200
201     ### First end-to-end A/B on the golden corpus (n=6) ? **ARCHIVED** *(2026-06-14)*
202     First real use of the harness + C1?C4 helpers, 100% offline (re-scored persisted WASB
trajectories, no GPU).
203     **Gate-A (net-crossings ? 2) is a measured win** ? F1 +0.074, 6/6 matches, sign-test
**p = 0.031**, CI
204     [+0.042, +0.104], over-segmentation (segment-count ratio) **1.68 ? 1.01**. The window-
level 5-feature logistic
205     **did not clear the bar** (?0.081, 2W/4L; a threshold/calibration artifact ? *not* the
stronger per-frame
206     prototype). The framework held up: **C1 promoted one variant and refused another; C4
surfaced over-segmentation
207     tIoU hid.** Full record + tables ?
208     [archives/quality-iterations/2026-06-first-ab-experiment/](archives/quality-
iterations/2026-06-first-ab-experiment/README.md).
209     Next: Gate-A is a candidate for the no-regression promotion gate; the logistic failure
is the C2 (calibration)
210     + threshold-in-loop task list.
211
212     ### Eval-honesty helpers (course corrections C1?C4) ? additive, default-OFF
*(2026-06-14)*
213     **Why.** A multi-agent literature review (folded into
214     [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) §13) confirmed the
framework's *skeleton* is
215     standard prior art (Temporal Action Localization/Segmentation + late/score-level fusion
? cite, don't claim),
216     and flagged that our **measurement** is walking into known small-N traps that would make
"promote only on
217     measured lift" promote noise or hide a real win. Four corrections, shipped as **pure,
additive helpers** (no
218     existing `compare`/LOVO/`metrics.score` behaviour changed ? callers opt in):
219     - **C1 ? never a single mean.** `backend/eval/eval_stats.py`:
`mean_with_ci`/`bootstrap_ci` (deterministic,
220     seeded), `paired_sign_test` (the exact distribution-free "won on enough *videos*?"
test the corpus already
221     needs), `paired_delta_ci` (B?A with a CI that must clear 0). At n?6 a lone LOVO mean
is not trustworthy.
222     - **C2 ? calibrate cue confidence.** `backend/eval/score_calibration.py`: Platt +
isotonic (PAVA) calibrators
223     + `brier_score`/`expected_calibration_error`, so a CNN softmax, an audio onset, and an
integer net-crossing
224     count are comparable before the arbiter and the over-confident cue can't swamp the
rest.
225     - **C3 ? honest stacking.** `eval_stats.out_of_fold_predictions`/`grouped_folds` (leave-
one-group-out by
226     match) so a *learned* producer's outputs used as arbiter features are out-of-fold (no
leakage).
227     - **C4 ? over-segmentation-visible metrics.** `backend/eval/segmentation_metrics.py`:
`f1_at_overlaps`
228     (F1@{10,25,50}), `average_precision`/`mean_average_precision` (mAP@tIoU),
`segment_count_ratio`. tIoU alone
229     is blind to a rally split in 3 or two rallies merged ? exactly what the reconciler
exists to fix.
230
231     **Status.** Helpers + 24 unit tests landed (pure stdlib, run in the numpy/GPU-free
lanes); ruff + mypy clean.
232     **Not yet wired into the default harness flow** (wiring `compare`/`run_regression.py` to
report CIs + F1@k is
233     the follow-on, owner-gated). The discipline they encode is the methodology for every
future cue measurement ?
234     see §13/C1?C4 for the citations and the remaining (Tier-2/3) corrections.
235
236     ### Multi-signal fusion P2: the FusionSegmenter *(in progress, 2026-06-14)*
237     **Hypothesis.** Adjudicating shuttle-seeded windows with the net-crossing **Gate A**
(`min_crossings?2`)

```

238 and a person-occupancy ****Gate B**** raises ****precision**** on the held-shuttle / warm-up /
one-sided-feed /
239 spectator false positives ? ***without*** hurting recall, because the shuttle stays the
primary,
240 camera-invariant seed.
241
242 ****Scope ?** and an honesty boundary from the archive (read this).****** The
243 [``2026-06-multivideo-lovo``](archives/quality-iterations/2026-06-multivideo-
lovo/README.md) iteration
244 already smoke-tested ****court-region** detection as a fix for the ***multi-court contrast***
confound****** and
245 found it ****low-headroom****: only 12?14% of dead-time shuttle detections fall ***outside***
the foreground
246 court, so a perfect court gate lifts contrast to only ~1.6?2.2× (below the ?3.6× "works"
threshold) ? the
247 multi-court problem is ****temporal****, and the ****learned model****, not a court mask, is its
answer. So Gate B
248 here is ****NOT**** sold as the multi-court fix. It targets a ****different, previously-**
untested****** failure mode
249 (held-shuttle/warm-up/one-sided/spectator FPs). Two distinct levers; don't conflate
them.
250
251 ****Data support for Gate A.**** The same archive shows real rallies on the multi-court
matches contain
252 ****?2 net-crossings 97% of the time (median 5)**** while a single service feed crosses
~once ? ``min_crossings?2``
253 is a principled, data-backed discriminator (it's the held-shuttle owner-bug killer).
254
255 ****Shipping shape.**** A NEW ``fusion`` segmenter, ****default OFF**** (``default_segmenter``
unchanged); Gate A
256 wired into the live trajectory path behind ``indexing.min_crossings`` (default 0 = OFF ?
wasb/tracknet
257 byte-identical today); the court mask is an ****optional**** config polygon (omit ? Gate B
is
258 player-count-only; supply ? masked). Person features (``players_in_court``,
``two_sided_motion``,
259 ``mean_player_motion``, ``in_court_ratio``, ``shuttle_player_colocation``) persisted into candidate
``stats`` for P3's
260 learned fusion. Reuses ``PersonDetector`` (P1) + ``rally_gate`` (already existed) ? no
reinvention.
261
262 ****Status (2026-06-14): SHIPPED, default-OFF.**** ``FusionSegmenter``
(``backend/pipeline/segmenters/fusion_hybrid.py``,
263 registry key ``fusion_hybrid``) + the pure feature math
(``backend/pipeline/detectors/fusion_features.py``) + two
264 identity hooks on the trajectory engine (``_enrich_candidate_stats``,
``_select_for_handoff``) + ``PersonDetector``.
265 ``detections_per_frame`` + the ``FusionConfig`` block. Built design-first (a
``fusion-p2-design`` fan-out workflow), then
266 ****adversarially reviewed by a 31-agent workflow (25 findings ? 19 confirmed)****: the
headline confirmation is
267 ****NO-REGRESSION verified**** ? wasb/tracknet twins persist byte-identical stats and hand
off all windows (pinned by
268 the extended equivalence test); default ``FusionSegmenter`` == its wasb substrate. The
confirmed low/medium findings
269 were fixed (frame-dim guards, ``round()`` frame indices, the per-window axis-estimate
redundancy ? a reusable
270 ``rally_gate`` estimate, honest torch-import docstrings, and 6 new tests closing the
person-detector / Gate-B /
271 substrate / colocation-halo gaps). Suite 441?467 green******; ruff + mypy clean.
272 ****Two findings deferred with intent**** (reviewer-agreed): person features are computed
eagerly even for
273 Gate-A-pruned windows (a P3 redesign, would change the persisted-superset semantics);
and ``fusion.velocity_threshold``
274 must be set equal to the substrate's to A/B against a tuned wasb run (documented in

```

`FusionConfig`).
275     **Lift still to be measured** offline on the 6 golden videos via
`experiment.compare(CONTROL, fusion)`; the
276     held-shuttle-drop + person-feature values are GPU/owner-side. *(Measured numbers
appended here when run ? no claims before.)*
277
278     ### Person cue extracted to a reusable producer (fusion P1) ?
[#115](https://github.com/avidullu/badminton-highlight-indexer/pull/115), 2026-06-13
279     Lifted the torchvision detector + ByteTrack + velocity windowing out of `yolo_hybrid`
into
280     `backend/pipeline/detectors/person_detector.py::PersonDetector`. **No behavior change**
(proven by the
281     unchanged `process_video` output tests); `yolo_hybrid` now orchestrates only. **Why it's
a quality move:**
282     it makes "add the person cue to fusion" an *add-a-producer*, not a segmenter rewrite.
Suite 440?441 green.
283
284     ### Experiment harness: A/B without rollback risk ?
[#114](https://github.com/avidullu/badminton-highlight-indexer/pull/114), 2026-06-13
285     Shipped `backend/eval/experiment.py` ? the thin variant/experiment layer (the audit's
"-80% already
286     exists" held; pure glue over
`metrics`/`calibration`/`classifier`/`training`/`regression.gt_hash`). Gives
287     `Variant` (+ pinned `CONTROL`), `compare()` (labeled A/B, LOVO-honest),
**`compare_unlabeled()`** (the SxS
288     read on NEW unlabeled footage ? agreed / A-only / B-only), and a JSON experiment
leaderboard. **This is
289     the instrument every iteration is measured on.** 19 tests incl. the no-regression
invariant (all cues OFF
290     ? byte-identical to control). Suite 421?440 green.
291
292     ---
293
294     ## Archived predecessors (the historical trail)
295     Completed iterations live under [`archives/quality-iterations/`](archives/quality-
iterations/README.md):
296     - **`2026-06-initial/`** ? seed of the quality phase.
297     - **`2026-06-human-gt-tuning/`** ? first eval vs **human** GT on a partial-label video;
partial-label
298     harness + labeled-window methodology; baseline 0.650 ? tuned 0.742 (*fit-on-self*);
found+fixed a
299     negative-start ingestion bug (collector PR #13).
300     - **`2026-06-multivideo-lovo/`** ? the **n=6 corpus + honest LOVO** ; established the
heuristic is
301     **footage-fragile** (multi-court Kushagra 0.16 / gBaaddy 0.28), the **"contrast"
metric** predicts
302     heuristic F1 monotonically, and a **learned per-frame prototype overtook the heuristic
at n=3 and holds
303     ~+0.10** (n=6: learned 0.583 vs heuristic 0.480), with **AUC 0.83?0.92 decoupled from
contrast**. This
304     is the data-backed green light for the owned model.
305
306     ---
307
308     ## Standing lessons (carry-forward)
309     - **Persist detection once, re-score variants forever.** Expensive/risky = per-frame
detection (GPU/WSL);
310     cheap/swappable = the decision. Most experiments never enter the served pipeline.
311     - **LOVO or it didn't happen.** Score on a held-out *video*, never a held-out window
from the same video.
312     Group by **match**, not file. Sign test needs n ? ~n=10/9-wins for p<0.05.
313     - **The multi-court confound is TEMPORAL, not spatial/loudness** ? three heuristics were
tested and failed
314     for the same reason (foreground spatial gate **?0.04** ; audio energy **AUC ~0.6** ;
court-region

```

315 separability ****12?14% out-of-region****): the dead-time confusion is rally-vs-non-rally
activity in the
316 ***same place and frequency band***. Only the learned model's temporal/structural features
crack it (AUC
317 0.83 on Kushagra/gBaaddy). ****Don't re-derive these dead ends.****
318 - ****Court-masking ? multi-court fix**** (see above) ? it's a ****precision**** tool for off-
court persons and
319 one-sided feeds, a ***different*** lever; measure its lift, don't assume it.
320 - ****Net-crossings are a strong rally signal**** ? 97% of real rallies cross ?2× (median
5); single feeds ~1×.
321 Gate A leans on this.
322 - ****Small-N feature discipline.**** Few, scale/translation-invariant features
(counts/ratios, not raw
323 coords) + a logistic regression, not a net ? raw coords memorize this court/camera.
324 - ****Calibration, not representation, is the small-N wall**** ? learned probabilities
separate (high AUC) but
325 the operating point doesn't transfer with few training videos. More distinct matches
is the highest-ROI
326 unblock.
327 - ****Gate A already existed, just wasn't wired to production**** ? the cheapest precision
win is plumbing, not
328 new capability. Always check what the eval drivers already do that the live path
doesn't.
329

14. user (2026-06-16T03:12:22.180Z)

1 # Next steps (live) ? owned-model + multisport
2
3 ****Updated:**** 2026-06-13 (****rally-detection quality track****: fusion + experiment-harness
design docs merged,
4 owner-gated ? see the dedicated section below, the desktop-agent pick-up). Prior:
2026-06-10 (audit?M0?Gemini-battery session; later same day the ****UI-alpha track opened****:
5 ADR-009 ratified ? KhelSutra Guru / khelsutra.guru, private `khelsutra` repo, local-
first delivery ? see
6 `docs/UI\_ALPHA\_EXECUTION\_PLAN.md`; ****PR-1 async jobs (#16) MERGED 2026-06-11, PR #87****;
****PR-2**
7 upload/download MERGED 2026-06-11, PR #99** ? live-E2E-verified (3-part chunked upload
MD5 byte-identity +
8 compile?browser-download); next: PR-3 identity. Same day: the ****8?9 quality sweep**
#90?#98 landed** ? ruff/mypy/
9 coverage-floor CI lanes, SQLite conn-close fix, config unknown-key warning, hybrid-twin
collapse into
10 `trajectory\_hybrid.py`, WSL tilde+abspath live-path fixes; suite 388 green; LOVO 0.626
re-verified exact).
11 ****Authoritative roadmap:****
12 `docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md` (+ ADR-008 for topology). ****Latest blow-by-**
blow:** memory
13 `current-state.md` / `session-handoff.md`. This file = the prioritized "what to do
next," refreshed each session.
14
15 **** ? Next-session / GPU-box pickup (2026-06-15)**
16 ****M0 in-container quick wins are COMPLETE:**** ship-gate provisional blocker (****#170****),
train/eval split guard (****#171****), the ****doc-consistency sweep**** (****#174**** ? corrected 16
drifted docs vs the code), and the ****P3 person-feature fusion litmus machinery** is already
present + suite-green** (`backend/eval/fusion\_compare.py` + `ablation.py`: person-driven ?F1 /
bootstrap CI / paired sign test; the real-golden `skipif` litmus waits only on GPU-extracted
data). ? CI is red repo-wide due to a ****GitHub Actions minutes/billing outage**** (account-level,
not code) ? work is local-gated + `--admin`-merged until Actions is restored (Settings ? Billing
? Actions minutes / spending limit).
17
18 ****Immediate next step ? the real P3 LOVO measurement, split GPU?CPU**** (the dev box is
CPU-only: `torch+cpu`/`cv2` present but no GPU, and the corpus/videos/CSVs/labels are not here):
19 1. ****[owner / GPU box]**** `python -m backend.eval.fusion_golden --manifest
output/human_lovo_manifest.json --out-dir "%USERPROFILE%\OneDrive\rally-annotator\golden-

```

data\features\2026-06-15-n6" ? produces `*_golden_features.json` for the 6 golden videos. (GPU
+ WSL + videos + trajectory CSVs + `rallies.csv` labels are owner-side only.)
20     2. Artifacts land in the shared cloud folder `C:\Users\avidu\OneDrive\rally-
annotator\golden-data\` (OneDrive auto-syncs both ways). Convention + workflow:
[[GOLDEN_DATA_SHARING.md]](GOLDEN_DATA_SHARING.md); tracked in ##175.
21     3. [CPU dev/agent session] `python -m backend.eval.fusion_compare --features-dir
<shared>\features\2026-06-15-n6 --record` + `python -m backend.eval.ablation --features-dir
<shared>\... --granularity group` ? real person-feature ?F1; copy the JSONs into `output/` to
also pass the `skipif` litmus tests (`tests/test_ablation.py::test_litmus_reproduces_gate_a`).
22
23     Also queued (owner-gated / pending input):
24     - Ship-gate target calibration ? ##172 (owner decision; needs corpus growth).
25     - Tests reorganization ? blocked on the owner's grouping metric (provide it ? an agent
will do it).
26     - Restore CI ? the Actions outage is account-level (billing/minutes), not a repo/config
bug.
27
28     Leaving (bulky / M2 ? no start without owner go): ActionFormer (M0.4), event-index
DB migration (M0.1), Gemini-silver purge (M0.3), cost-to-Gen-2 benchmark + ByteTrack
deprecation, multisport, repo rename, PMF execution.
29
30     Where we are (one paragraph)
31     ADR-008 ratified ("repo split driven by productionization, not isolation"): training
lives in-repo (`training/`)
32     behind a CI-enforced import wall; the featurizer is single-sourced
(`backend/features/rally_seq.py`, train==serve);
33     the Gen-0 harness ships (`python -m training.gen0.harness --manifest
output/human_lovo_manifest.json --floor
34     eval_baselines/heuristic_lovo_n6.json --version <semver>`) and produced artifact v0.1.0
(LOVO 0.626, floor passed,
35     sign test 5/6 wins p=0.109 ? honestly not significant, ship=False). The n=6 owned
corpus is in the collector
36     (262 rallies, Proprietary, commercial export = owned data only after the ShuttleSet
reclassification). The Gemini
37     battery is done (see table) ? the moat is quantified.
38
39     The quality table that now drives strategy (IoU 0.5 vs human golden labels)
40     | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |
41     |---|---|---|
42     | Heuristic (velocity windowing) | 0.657 | 0.157 |
43     | Gen-0 owned head (LOVO, n=6) | 0.744 | 0.561 |
44     | Two-stage WASB+Gemini refine | 0.83 | ? |
45     | Gemini wrapper (`gemini-flash-latest`) | 0.93 | 0.66 |
46
47     Reading: clean footage is commoditized (serve it with Gemini for cents). Multi-
court drops the commodity
48     wrapper to 0.66 ? that is the owned model's ship target (its FPs are background-game
pickups + dead-time merges,
49     the exact contrast-metric confound). Gen-0 is 0.10 behind with only n=6 ? a corpus-
growth-sized gap, not an
50     architecture-sized one. Two-stage refine = the cost-efficient serving path for LONG
tapes (uploads candidate clips
51     only); wrapper 503-flakiness (observed all session) makes the provider-fallback registry
load-bearing.
52     Baselines tracked in `eval_baselines/` (heuristic_lovo_n6.json,
gemini_wrapper_2026-06-10.json).
53
54     Rally-detection quality track (design merged 2026-06-13, owner-gated; implementation
in progress)
55     This is the track a desktop/local agent is slated to pick up. Design docs landed (PR
#112 + session). All implementation is owner-gated (per explicit in the plan); each step =
ONE PR off `master`.
56
57     Hardening loop (T1-T5 positive tests + audit + finale) COMPLETE (2026-06-14, PRs
#156/#157/#160/#161/#163 + #165 + final records). See archived

```

```

`docs/archives/CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md` (Living Log with every
pre/post/review) and `docs/archives/HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md` (full rules
+ STATE). Top-level now pointer stubs per #165 review. All 5 items + verification + archive
done; clean slate.
58
59 - **Docs:** `docs/MULTI_SIGNAL_FUSION_PLAN.md` (shuttle + person + audio, one fusion
layer) .
60     `docs/EXPERIMENT_HARNESS.md` (A/B + shadow + promotion gate, zero-regression) .
61     `docs/ROORA_LABELING_GUIDELINES.md` (v8) . `docs/HOW_RALLY_DETECTION_WORKS.md`
(2-layer mental model).
62 - **Key finding (no re-investigation needed):** the owner's "held-shuttle counted as a
rally" bug is **NOT a missing capability** ? `backend/pipeline/detectors/rally_gate.py` already
implements the net-crossing "Gate A" (`apply_gate(..., min_crossings=2)`) that kills
service-feed/held-shuttle windows, but it is only called from eval drivers (and there was no
`min_crossings` knob in `IndexingConfig`). Wiring Gate A into the live segmenter (now via
fusion_hybrid) was the single cheapest, highest-confidence quality win. (P2 FusionSegmenter
specifics: registry `fusion_hybrid`, default-OFF, seeds from substrate, persists
`net_crossings` + optional person features, optional served Gate A/B via config knobs;
adversarially reviewed; suite green. See `tests/test_served_gate_a_regression.py` for the honest
finding that on production windowing Gate A is neutral/negative in some cases, hence kept opt-in
with knob for safety, and the leaderboard revert note.)
63 - Current status (2026-06-14): P1 (person cue extraction to reusable
`PersonDetector`) + P2 (FusionSegmenter + Gate A signal + served handoff gating) + experiment
harness P1 DONE (adversarial review + NO-REGRESSION; suite green). P3 next (this PR + follow-
ups): learned fusion (person features ? harness `compare` LOVO lift on the 6 golden videos;
eager-person-compute optimisation); then P4 audio via hit_detector.
64 - Pick-up order (each = ONE PR off `master`, owner approval first, flag-gated default-
OFF, promote only on measured LOVO):
65     1. (P2 first-step, largely done via fusion_hybrid) Wire Gate A into production served
path + config knob (min_crossings); apply in hybrids at runtime. (Now available opt-in; served
neutral on production windowing per litmus ? kept opt-in with knob for safety.)
66     2. P3 first step (this work item): harness `compare` LOVO litmus for person-on variant
(players_in_court + two_sided + colocation etc. from `fusion_features`) on the 6 golden; add the
eager compute optimisation note/landing. Pure + harness (testable here).
67     3. P3/P4 continuation: audio cue integration + full learned promotion if lift proven.
68 - Discipline (from EXPERIMENT_HARNESS.md): every new cue flag-gated, default
OFF; promote to default only on measured LOVO lift through `run_regression.py` (exit
0/1/3); pinned `CONTROL` ? rollback = config flip. Harness re-uses ~80% existing code.
69 - ? Dev container no GPU/torch/cv2 ? Gate A, person cue logic, harness, and pure fusion
paths are testable here; real-video / golden LOVO / ablation confirmation is owner-side (GPU/WSL
machine).
70
71 Discipline reminder: owner approval before starting any owner-gated item; one PR per
step; babysit the PR (poll/fix/reply) until merge observed, then advance.
72
73 ## Prioritized next steps
74 1. [owner, in progress] GROW THE GOLDEN CORPUS ? the single highest-leverage move;
every lever below feeds on it.
75     Priority order for new videos: (a) multi-court badminton (the target distribution
AND where the ship target
76     lives), (b) ?3 matches per new sport ? table tennis first (WASB transfers at F1
0.909; pickleball labels are
77     corpus-gold but not trainable until a clean MIT/Apache ball detector exists), (c)
capture variety per
78     `VIDEO_RECORDING_GUIDELINES.md`; adult sessions only for the training pool
(consent guardrail). Per video:
79     label ? `curate import-local --normalize --license Proprietary --fps <true fps>` ?
re-run the Gen-0 harness.
80     The paired sign test needs n: at n=10 with 9 wins, p?0.011 ? significance is
reachable this month.
81 2. [owner decision] Set the ship-gate target ? now informed: floor = heuristic 0.480
(necessary), multi-court
82     reference = wrapper 0.66 (the number that makes the owned model worth serving),
clean-footage ceiling = 0.93
83     (not the target ? Gemini serves that tier). Suggested shape: "LOVO F1@tIoU0.5 ? 0.70

```

on multi-court folds with

```

84     paired-sign p<0.05 vs heuristic" ? owner to ratify/adjust.
85     3. **M0.4 next increment ($0/local):** swap the Gen-0 LogReg for
**ActionFormer/ASFormer** (MIT, minutes on the
86     2070S ? still lightweight, NOT the heavyweight Gen-2 stack) inside the existing
harness; per-video threshold
87     calibration (the LOVO-vs-oracle gap is visible per fold in the harness output).
88     4. **Multisport thread:** ingest Roora pickleball/TT CSVs (`import-local --normalize`);
merged-prototype LOVO across
89     badminton+TT once TT has ?3 matches; identify a clean pickleball ball detector. ? TT-
on-WASB serving stays gated
90     by C3 (provenance multiplies per sport).
91     5. **PMF validation (cheaper than any engineering month):** C13 academy interviews (+
the two added questions:
92     "multi-court footage nobody watches?" / "pay per-match to index it?"); camera pilot
at the warmest 2-3 academies
93     (consent at capture; demo served via the Gemini path or human-polish ? the reel
pattern exists:
94     `output/MahadevpuraSingles_highlights.mp4`). **Field kit COMPLETE (2026-06-12, owner
ratified the
95     physical-visit approach; reframed product-feedback-first per owner spec):**
96     `docs/PRODUCT_FIELD_ASSOCIATE_FLYER.md` (the recruiting flyer ? owner fills ?/contact
placeholders and recruits) +
97     `docs/FIELD_VISIT_PLAYBOOK.md` (the hired associate's visit playbook: coach + student
script, demo-after-Q6
98     rule, ground rules, Phase-2 record-and-show-back stretch goal) +
`docs/FIELD_VISIT_ANSWER_SHEET.md` (per-visit
99     capture form incl. demo-reaction + student blocks) + `docs/PMF_FIELD_SIGNALS.md`
(G/A/R anchors + **pre-registered
100    decision rules** ? PROCEED / PIVOT A-C / FALSIFIED thresholds fixed before the first
visit).
101    6. **Housekeeping:** ~~rotate the Gemini API key~~ **DONE 2026-06-10** (owner-verified
working; confirm old-key
102    deletion in AI Studio); **rename repo off "badminton"** (owner action,
103    in BACKLOG); collector=SoR confirmation; off-peak niceties: none pending ? the
battery completed.
104
105    ## Pointers / artifacts
106    - Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` . Topology: ADR-008
(`docs/archives/decisions/DECISIONS.md`) .
107    Study: `docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md` . Quality history:
108    `docs/archives/quality-iterations/2026-06-multivideo-lovo/`.
109    - Harness: `training/gen0/harness.py` (one command ? train ? LOVO ? gate ? artifact
bundle; weights gitignored,
110    manifest reviewable). Featurizer: `backend/features/rally_seq.py` (FEATURE_SCHEMA v1).
111    - Corpus: collector DB (system-of-record, 6 matches/262 rallies) . eval manifest
`output/human_lovo_manifest.json`
112    (gitignored) . durable trajectories + Roora files: `C:\Users\avidu\Projects\Annotation
Setup\Collect\Trajectories\` .
113    labels: `...\Collect\Golden Labelled\`.
114    - Gemini ops learnings: serial uploads (`indexing.ai_max_workers=1`) + chunk re-encode
115    (`indexing.ai_chunk_scale_width=1280`) + 6-attempt exponential generate retry ? a
cold/prepaid project burst-throttles
116    with a MISLEADING "credits depleted" 429; ramp gradually.
117    - ? No owned-model implementation beyond greenlit M0.x until owner greenlights the next
milestone. Committed next
118    PLATFORM slice = async job model (single-tenant).
119

```

15. user (2026-06-16T03:12:22.648Z)

```

1     # Personalization Plan ? a per-user feature on a generalization-first baseline
2
3     **Status: DESIGN ? owner-adopted (hybrid); Phase-0 engineering landing.** This is the
authoritative

```



```

4      plan for the personalization track. The detailed strategy analysis + the full pros/cons
that led here
5      live in the working artifacts `scratch/PERSONALIZATION_PIVOT_ANALYSIS.md` (the decision-
grade
6      counter-case) and `scratch/PERSONALIZATION_FEATURE_DESIGN.md` (the code-grounded
design); this doc is
7      the checked-in summary + decisions + phased build plan. Peer to
8      [OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md),
9      [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md), and
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md).
10
11     ## 0. Thesis & strategic rationale (why *hybrid*, not personalization-first)
12
13     **Foundation principle (owner):** "a tool that gets better the more you help it."
Contribution
14     (data / video / diagnostics / annotation) improves BOTH the user's own model AND the
shared baseline;
15     **payoff scales with contribution.**
16
17     A directional pivot ? *personalization-first* (ship a baseline that continually self-
optimizes to one
18     user's own footage) ? was researched in depth. The verdict was **HYBRID, not all-in**:
19
20     - **Adopt personalization as a flag-gated FEATURE on a generalization-first baseline.**
The harness
21     genuinely makes per-user *tuning* near-free (it is already a corpus-relative A/B
engine; a per-user
22     corpus is just a special case). This is real and worth shipping.
23     - **Do NOT bet the strategy on per-user *learning*.** Per-user A/B is the n?6 small-
sample regime
24     *worse* (sharding data when you should pool); the differentiated parts (a learned per-
user model
25     serving path, the correction?label write path, a user/tenant dimension) are largely
unbuilt; a
26     naively-overfit per-user model loses the cross-user data flywheel and has no shared
baseline to
27     regress against.
28     - **The owner's contribution-flywheel + free-tier-pooling closes the one real hole** the
counter-case
29     found: by having (free-tier) users contribute to the *shared* baseline, the
compounding cross-user
30     moat is preserved *alongside* personalization ? personalization **and** flywheel, not
either/or.
31
32     Net: a generalization-first **baseline** ("batteries"), a thin **per-user tuning** layer
gated by the
33     honest small-N statistics (with a hard floor + structural-guard protection), and a
**contribution
34     loop** that feeds both.
35
36     ## 1. Owner decisions (locked 2026-06-14)
37
38     1. **Identity** ? auth WILL land; *design for the north star (multi-user / cloud),
implement for
39     current use (single-user local)*. `user_id` nullable now (NULL = local install),
device/install
40     UUID for the alpha; real auth + a `users` table later.
41     2. **Per-user model storage** ? inline `classifier_json` in the (local) DB now, but
**behind a load
42     interface** so that post-auth it is fetched from the **cloud-hosted DB run by the
service** that
43     generates base models and orchestrates the product. The local store is one
implementation of that
44     interface.
45     3. **`min_n` is a PROMOTION-confidence floor, NOT a service/segmentation gate.** A

```

```

brand-new user with
46     one video is ALWAYS fully segmented by the baseline batteries ? never denied, never
"give me more
47     videos first." `min_n` only gates whether a *per-user* A/B is trusted enough to
OVERRIDE a baseline
48     cue default *for that user*; below it (or if the CI / sign test don't clear) ? keep
the baseline.
49     **Training-corpus policy:** *no more than needed, no less than all-if-few* ? fit on
the
50     `smallest_k_with_lift` videos (the held-out-USER learning curve), else ALL available.
51     4. **"Flip Gate-A to default"*** = ship the validated Gate-A held-shuttle guard as the
**default
52     detection path** (today it lives in the opt-in `FusionSegmenter`, which is not the
default
53     segmenter). Changes served behavior ? **gated on a golden-set regression run + owner
OK**; still
54     open (see §5).
55     5. **Flywheel = backend property** (not in-product marketing). A **free tier** is likely
and
56     **free-tier videos ARE pooled** into the baseline (consent baked into the free terms;
57     minors-exclusion still enforced) ? the clean value-exchange that feeds the owner-
trained base model.
58     Paid tiers get stronger privacy / opt-out.
59
60     **Pending sub-decision:** structural-guard scope ? recommended **Gate-A only** (the
proven structural
61     held-shuttle guard; its value is in failure modes a user's corpus may not yet contain),
with **Gate-B**
62     (person-occupancy / `min_players`) left *subject* to per-user A/B; **`min_n = 5`**.
**Gate-A-only
63     structural, `min_n = 5` ? owner-confirmed 2026-06-14.**
64
65     ## 2. Architecture
66
67     **Two-layer decoupling (already real).** Detection (GPU, run once): `FusionSegmenter`
persists the
68     **full superset** of per-window features into `candidate_segments.stats`; only the
served handoff is
69     gated. Decision (CPU, swappable): `experiment.Variant` + `predict()` re-score persisted
features.
70     **Personalization == a per-user `Variant`** (cue_flags + `min_crossings` / `min_players` +
optional
71     LogReg + threshold). No segmenter rewrite ? a per-user `FusionConfig` overlay resolved
upstream.
72
73     **The single highest-leverage wiring:** `fusion_hybrid.py::_select_for_handoff` is the
production
74     decision seam. When a per-user model is pinned AND the promotion gate says ON, apply
75     `LogisticRegression.predict_proba` over `window_stats` (the same math
`experiment.predict()` runs
76     offline). The **load half of the bridge already exists** ?
77     `backend/pipeline/personalization/serving.py::load_user_classifier` rehydrates the
stored
78     `classifier_json` (and `resolve_user_fusion_config` overlays the per-user
`fusion_config`) from
79     the v5 store ? but it is **not yet wired into the served path**: `_select_for_handoff`
still reads
80     only the static `idx_cfg["fusion"]`. Closing that one seam (the first change to served
behavior, so
81     owner-gated) is the remaining #1 wall.
82
83     **Per-user data model** (migrations follow the `PRAGMA user_version` ladder; current =
v5):
84     - **v4** ? ? **shipped.** Nullable `user_id` on `videos` + `candidate_segments` (NULL =
local;

```

```

85     byte-identical to today).
86     - v5 ? ? shipped. user_models store: `(user_id, version, baseline_version,
87         feature_schema_hash, fusion_config, classifier_json, promotion_evidence,
n_videos_at_fit,
88         is_active)`. Inline classifier_json (tenant-safe); feature_schema_hash is the
89         MINOR-vs-MAJOR upgrade switch.
90     - v6 ? ? pending. contribution_consent ledger (default-deny; minors-exclusion
enforced at
91         the pooling query, not advisory).
92
93     ## 3. Constructs (the gap list + what is built)
94
95     | # | Construct | Status |
96     |---|---|---|
97     | 1 | Per-user store + serving bridge ? store + load half built (v5 user_models +
serving.py::load_user_classifier); only wiring _select_for_handoff remains | Phase 1 (the #1
wall ? load half done, wiring pending) |
98     | 2 | Correction-loop ? per-user labels (human_label setter + confirm/reject API;
recall-protecting de-bias prior) | Phase 2 (human_label is a read-only stub today) |
99     | 3 | Per-user promotion gate (min-N floor + structural-guard exemption) | ? DONE
? PR #141 (backend/eval/per_user_gate.py) |
100    | 4 | Contribution flywheel (opt-in data/video/diagnostics ? shared baseline; payoff ?
contribution; consent + minors-exclusion) | Phase 2 |
101    | 5 | User-as-annotator charter (rally-annotator: auto-generate a ~2-min high-
uncertainty segment from the user's own video; guide labeling ? grounding) | Phase 3 |
102    | 6a | Held-out-USER learning curve (the falsifiable "<K videos to superior quality"
criterion) | ? DONE ? PR #142 (backend/eval/per_user_eval.py) |
103    | 6b | Per-user no-regression guard (frozen per-user CONTROL; never accept a regressing
retrain/upgrade) | ? in flight (feat/per-user-no-regression) |
104    | 7 | Day-0 "batteries" = the validated label-free Gate-A/B fusion (?F1 +0.074, 6/6,
p=0.031) shipped as the default | Phase 0, owner-gated ($5) |
105
106    ## 4. Phased PR-sized build plan + status
107
108    Phase 0 ? safety + day-0 batteries (no schema / API / served-behavior change; pure-
Python + tests):
109    - PR-0.1 per-user promotion gate ? ? #141 merged
110    - PR-0.2 held-out-USER learning curve ? ? #142 merged
111    - PR-0.3 per-user no-regression guard ? ? in flight
112    - PR-0.4 ship Gate-A as the day-0 battery ? owner-gated ($5; needs a golden-set
regression run)
113    - PR-0.5 within-user LOVO measurement (the "confirm-or-kill" experiment) ? owner-
gated (needs multi-video-per-user data on the GPU box)
114
115    Phase 1 ? per-user tuning serving + store ? already landed ? v4 migration (nullable
user_id),
116    v5 user_models + CRUD (upsert_user_model / get_active_user_model), and
serving.py
117    (resolve_user_fusion_config overlay resolver + load_user_classifier). Remaining:
set_user_id
118    (and the Phase-2 set_human_label writer); wire the bridge into _select_for_handoff
(owner-gated
119    ? served path); /retune + /model trust-feature endpoints.
120
121    Phase 2 ? correction loop + flywheel: POST /api/candidates/{id}/label (the first
human_label
122    writer); recall-protecting de-bias prior; v6 consent ledger + minors-exclusion; opt-in
contribution
123    upload (owner-gated ? privacy surface); free-tier pooling.
124
125    Phase 3 ? annotator charter + taxonomy + baseline upgrades: active-learning task
picker; rally-
126    annotator CSV ingest ? per-user labels; coarse hard/soft shot proxy (owner-gated);
baseline-
127    upgrade governance (MINOR carry / MAJOR re-fit on feature_schema_hash) + optional

```

```

federated nudge
128     (owner-gated, counsel).**
129
130     ## 5. Open decisions (owner)
131
132     1. **Structural-guard scope + `min_n` value** ? lock **Gate-A-only structural, `min_n =
5`** (recommended)?
133     2. **Gate-A default flip (PR-0.4)** ? Gate-A lives in the non-default `FusionSegmenter`,
so "ship as
134     default" is a path choice (make that segmenter the default vs port the gate into the
default path).
135     Recommended sequence: **run a golden-set regression first**, bring the number, *then*
approve.
136     3. **Per-user `user_id` for the alpha** ? device/install UUID acceptable, or must auth
land first?
137     4. **Local-install flywheel default** ? fully-local zero-pooling until consent is
explicitly granted?
138
139     ## 6. Honest risks (carried from the analysis)
140
141     - **Small-N per-user eval noise** is central and *worse* per-user ? mitigated
structurally: the gate
142     refuses below `min_n` and falls back to the baseline; never auto-promotes in the
regime the tooling
143     exists to refuse in. PR-0.5 confirms-or-kills the core premise before any serving
build-out.
144     - **Cold-start** ? the day-0 Gate-A battery (zero per-user labels) is the floor for the
median user.
145     - **Privacy deferral** ? do not advertise "data never leaves your machine" until on-
device ships;
146     near-term pooling carries the full cloud (DPA / consent / minors) posture.
147     - **Compute** ? heavyweight VLM personalization (on-device Qwen3-VL fine-tune) is **out
of v1 scope**;
148     v1 is the cheap LogReg / threshold / cue-ON-OFF layer only.
149
150     ## 7. References
151     - `scratch/PERSONALIZATION_PIVOT_ANALYSIS.md` ? the decision-grade strategy analysis
(pros/cons, harness-fit, counter-case, recommendation).
152     - `scratch/PERSONALIZATION_FEATURE_DESIGN.md` ? the code-grounded design (schema,
serving bridge, APIs, critical files).
153     - Built: `backend/eval/per_user_gate.py` (#141), `backend/eval/per_user_eval.py` (#142).
154     - Harness: [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md), `backend/eval/experiment.py`,
`backend/eval/eval_stats.py`.
155

```

16. user (2026-06-16T03:12:22.968Z)

```

1     # Roora Annotation Guidelines ? Rally Golden Set
2
3     **Version: v8 (2026-06-13).** Supersedes the older Drive copies (v3/v5/v6/v7) **and**
the internal
4     "DRAFT v2" label ? those numbers never matched each other; **v8 is the single label to
use with
5     Roora going forward.** Content is otherwise the finalized 2026-06-12 guideline plus the
6     fusion-driven additions marked **`? ADDED (fusion)`** below.
7
8     > **?? Canonical home / provenance.** **This in-repo file is the sole canonical copy of
the guideline**
9     > ? it lives here because the multi-signal-fusion features that motivate the new fields
(court
10    > calibration, singles/doubles) originate in *this* repo (see
11    > [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) and
12    > [GOLDEN_SET_VENDORIZING.md](GOLDEN_SET_VENDORIZING.md), ADR-006). The annotation
deliverable still imports
13    > into the sibling **`sports-data-collector`** repo via `python -m src.cli.curate

```

```

import-local`; if a
14     > mirror is ever staged there for the collector's intake docs, it must be regenerated
from this file
15     > (and the `? ADDED (fusion)` callouts are exactly the deltas to port).
16     >
17     > TODO(avidullu)` markers are preserved` ? they are owner-only items (Drive links,
intake
18     > manifest, NDA counsel review) and must be filled before sending to Roora. Search for
`TODO(avidullu)`.
19
20     Brief (vendor spec) + Illustrated Guide (rater how-to), in one doc.
21
22     ---
23
24     # PART A ? Roora Annotation Brief (vendor-facing spec)
25
26     Project: racket-sport rally segmentation golden / regression dataset.
27     Sports in scope: Badminton, Pickleball, Table Tennis (each in singles AND doubles).
28
29     This brief is the vendor-facing spec. The annotation deliverable is a per-rally CSV that
imports
30     directly via `python -m src.cli.curate import-local`.
31
32     ## 0. What we need in one sentence
33
34     For each provided match video, produce ONE CSV file listing EVERY rally with its start
time, end
35     time, number of shots, and how it ended ? EXCLUDING all dead time.
36
37     ## 1. Deliverable format (canonical)
38
39     One CSV per video (UTF-8), named exactly `.csv` (we supply the `video_id`).
One row per
40     rally, chronological, header row required.
41
42     Columns:
43
44     | Column | Type / values | Meaning |
45     |---|---|---|
46     | `rally_number` | integer, 1-based, chronological ? no gaps, no repeats | rally index |
47     | `start_mmss` | `m:ss.mmm` (e.g. `1:12.40`) | rater-entered serve-contact time |
48     | `end_mmss` | `m:ss.mmm` | rater-entered ball/shuttle-dead time |
49     | `start_time` | decimal seconds (e.g. `72.40`) ? AUTO-FILLED from `start_mmss`
| machine time |
50     | `end_time` | decimal seconds ? AUTO-FILLED from `end_mmss` | machine time |
51     | `shots_count` | integer | total racket/paddle/bat contacts (serve = shot 1) |
52     | `ending_reason` | `winner` \| `forced_error` \| `unforced_error` \| `service_fault` \|
`let` \| `other` | WHY it ended |
53     | `last_shot_outcome` | `in` \| `net` \| `out` \| `n_a` | WHERE the last shot landed
(`n_a` for service_fault/let) |
54     | `score_before` | optional, `"A-B"` ? blank if not confidently readable | score before
|
55     | `score_after` | optional, `"A-B"` | score after |
56     | `notes` | optional free text ? REQUIRED when `ending_reason = other` | edge-case
note |
57
58     Times (mm:ss ? decimal seconds): raters type only `start_mmss` / `end_mmss`. The
template Google
59     Sheet fills `start_time` / `end_time` with this formula (shown for `start_time`, row 2):
60
61     ```
62     =IF(B2="", "", IFERROR(VALUE(LEFT(B2,FIND(":",B2)-1))*60 +
VALUE(MID(B2,FIND(":",B2)+1,20)), VALUE(B2)))
63     ```
64

```

```

65     On "Download as CSV" the computed decimal seconds export. (If your tool already shows
decimal
66     seconds, type those into `start_time` / `end_time` and leave the mm:ss columns blank.)
67
68     **Example (badminton):**
69
70     ```
71     rally_number,start_mmss,end_mmss,start_time,end_time,shots_count,ending_reason,last_shot
_outcome,score_before,score_after,notes
72     1,0:12.50,0:25.00,12.50,25.00,9,winner,in,0-0,1-0,
73     2,0:45.00,0:47.30,45.00,47.30,1,service_fault,n_a,1-0,1-1,serve into net
74     3,1:01.20,1:18.85,61.20,78.85,14,unforced_error,out,1-1,2-1,final clear sailed long (no
pressure)
75     ```
76
77     **Critical format rules:**
78     - `start_time` / `end_time` that reach us MUST be decimal seconds (the formula
guarantees this).
79     - Times are relative to the **START OF THE FILE WE PROVIDE**. Do NOT trim, clip, or re-
encode the
80     video ? our pipeline identifies a video by its file checksum. (The file we share may
be an
81     **annotation-optimized copy** of the original ? lighter and faster to frame-step, but
82     frame-for-frame identical. Treat it exactly like the original: annotate it as-is.)
83     - `end_time > start_time` for every row.
84     - No overlaps: `end_time[N] <= start_time[N+1]`.
85     - Everything between one rally's `end_time` and the next rally's `start_time` is DEAD
TIME and gets
86     NO row.
87
88     ## 2. The three non-negotiable rules (full detail + pictures in the Guide)
89
90     1. **IGNORE ALL DEAD TIME.** Only label active play.
91     2. **INCLUDE EVERY RALLY ? EXCLUDE NONE.** Every served point, including service faults,
1?? shot
92     rallies, and lets/replays. A missed rally is the worst error in this project.
93     - **EXCEPTION:** warm-up / knock-up is NOT a rally ? pre-match and between-games
warm-up looks like
94     rallying but isn't scored play, so don't label it. Tell it apart by: no real serve;
cooperative
95     (not competitive) hitting; the score not progressing; often several shuttles/balls
in use (full
96     how-to in the Guide, Part 2). When unsure if the match has started, flag it.
97     - **SECOND EXCEPTION:** a rally already in progress when the file starts (or still in
progress when
98     it ends) gets NO row ? its true start (or end) is outside the footage and must
never be guessed
99     or extrapolated. Note it in the delivery instead.
100    3. **A SHOT = ONE CONTACT** with the racket/paddle/bat. Serve = shot 1. Bounces and net-
clips are not shots.
101
102    ## 3. Scope & phasing
103
104    - **Phase A ? Pilot (this engagement):** 3 videos (1 badminton + 1 pickleball + 1 table
tennis), each
105    < 3 min with ~10?12 rallies. This is the paid calibration round; we review jointly and
lock the
106    Guide before scale-up.
107    - **Phase B ? Scale (after a successful pilot):** ~12?15 videos PER SPORT (~36?45
total), each ~30 min.
108    Same format, rules, and template.
109
110    Doubles are included across the program. Rally rules are identical to singles; the
badminton- and
111    pickleball-specific doubles notes are in the Guide.

```

```

112
113     ## 4. Quality bar (acceptance criteria)
114
115     Automated validation (no overlaps, positive durations, schema-valid) runs on every file,
then a ~20%
116     human spot-check. A file is returned for rework (in scope) if it misses:
117
118     - Completeness: every rally present; zero missed rallies; zero dead-time rows.
119     - Boundaries: `start_time` / `end_time` within  $\pm 0.3$  s of true serve-contact /
landing.
120     - ?? Boundaries that all sit on a frame-sampling grid (e.g. every 30th frame)
STRUCTURALLY FAIL this
121     bar and are detected automatically ? the whole file is returned for rework. See §7
for the
122     required method.
123     - shots_count: exact for rallies ? 15 shots;  $\pm 1$  allowed for longer/faster rallies.
124     - ending_reason: one of the six allowed values, per the Guide's decision tree.
125     - Format: decimal seconds present, valid columns, non-overlapping, chronological.
126
127     > ADDED (fusion) ? court calibration as a day-1 quality lever. The owner's
pipeline produced a
128     > bad highlight where a held shuttle was mislabeled as a rally. The boundary definition
($ Part 3:
129     > start = serve contact, not the toss/hold) already excludes that ? so faithful
boundaries are
130     > themselves a quality checkpoint. The one NEW field we need (next section) is court
calibration,
131     > which lets the detector reason about court occupancy and net-crossings.
132
133     ## A. Court calibration & intake manifest ? ADDED (fusion)
134
135     > ADDED (fusion). Why: the multi-signal detector
136     > ([MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md)) needs to know where the
court is and
137     > where the net is to (a) ignore people outside the court, (b) require ?2 players on
court for a
138     > rally, and (c) count shuttle net-crossings. This is a one-time, per-video
annotation ? cheap,
139     > high-leverage, and required in deliverables from day 1.
140
141     Per-video, recorded once in the intake manifest (NOT in the per-rally CSV):
142
143     | Manifest column | Type | Meaning |
144     |---|---|---|
145     | `video_id` | string | (existing) we supply it |
146     | `fps` | number | (existing) we supply it |
147     | `court_corners` | 4 points (x,y) in pixels, clockwise from the near-left corner |
the playable court rectangle |
148     | `net_line` | 2 endpoints (x,y) in pixels | the net, left post ? right post |
149     | `singles_or_doubles` | `singles` \| `doubles` | sets the expected on-court player
count (2 vs 4) |
150
151     - Pick one clear frame with the empty/visible court and click the 4 corners + 2 net
endpoints.
152     - If the camera does not move for the whole video (the normal case), one calibration
covers the file.
153     - If the camera is moved mid-file, flag it ? we'll split the file.
154     - Worked example (badminton, 1920x1080):
`court_corners`=[(360,980),(1560,980),(1330,300),(590,300)]`,
155     `net_line`=[(470,640),(1450,640)]`, `singles_or_doubles`=singles`.
156
157     (After the first batch we may add more per-shot fields ? see the §B appendix ? and
renegotiate price.)
158
159     ## 5. Per-sport quick reference (full version + examples in the Guide)

```

```

160
161 - **Badminton (singles & doubles)** ? start: racket?shuttle serve contact . end: shuttle
floors /
162     fault . shot: each racket?shuttle contact.
163 - **Pickleball (singles & doubles)** ? start: paddle?ball serve contact . end: double
bounce / out /
164     net / fault . shot: each paddle?ball contact (floor bounces are NOT shots; two-bounce
rule is normal play).
165 - **Table tennis** ? start: racket?ball serve strike (not the toss) . end: failed legal
return . shot:
166     each racket?ball contact (table bounces are NOT shots). Net-cord serve = let.
167
168 ## 6. `ending_reason` values (decision tree in the Guide)
169
170 This vocabulary attributes WHY the point ended; it is shared across all net-separated
sports
171 (badminton, pickleball, table tennis, tennis). Record WHERE the last shot physically
landed in the
172 separate `last_shot_outcome` column (`in` | `net` | `out` | `n_a`) ? e.g. a forced error
that flew
173 long is `ending_reason=forced_error, last_shot_outcome=out`.
174
175 - `winner` ? last shot landed IN and the opponent made no legal contact (a clean
winner).
176 - `forced_error` ? the losing player erred (netted / out / missed) UNDER PRESSURE ?
stretched,
177     wrong-footed, or rushed by the opponent's strong shot.
178 - `unforced_error` ? the losing player erred on a ROUTINE, UNPRESSURED ball they had
time and position
179     to make.
180 - `service_fault` ? the serve itself failed / was illegal.
181 - `let` ? the point was replayed.
182 - `other` ? anything else / cannot tell. MUST add a note.
183
184 ## 7. Capture method (required for boundary frames)
185
186 Use any tool that shows current time (`mm:ss.mmm` or seconds, or frame + known fps) and
can step one
187     frame forward/back: a frame-steppable player (VLC frame-step) or a timeline tool (Label
Studio /
188     CVAT). The deliverable is the CSV in the schema above ? or, if you work in CVAT, the
CVAT export as-is
189     (we recompute each rally's time from its box frame range; seconds = `frame_number /
fps`, fps provided
190     per video).
191
192 **Boundary placement is the non-negotiable part:**
193 - Set up the CVAT task at **FRAME STEP 1** on the file we provide (it is encoded to make
this fast ?
194     frame-stepping and back-scrubbing stay snappy).
195 - The FIRST box of every rally goes on the EXACT serve-contact frame; the LAST box on
the EXACT
196     shuttle/ball-dead frame. Frame-step to find them.
197 - Boxes IN BETWEEN may stay sparse (e.g. every 30th frame) ? only the two boundary
frames must be exact.
198 - Self-check before submitting: every start/end within  $\pm 0.3$  s of the true frame (Guide).
199
200 ## 8. Logistics & contacts
201
202 - **TODO(avidullu):** paste the Google Drive folder link where source videos are shared
with Roorā.
203 - Per-video fps + `video_id` (+ the calibration columns from SA): provided in the intake
manifest
204     (template `INTAKE_MANIFEST_TEMPLATE.csv`). **TODO(avidullu):** fill the manifest (one
row per video)

```



```

205     and share its link.
206     - **TODO(avidullu):** paste the Drive folder link where Roora uploads completed
    `<video_id>.csv` files.
207     - Blank annotation template Sheet (mm:ss ? seconds formula from $1 pre-filled):
208     **TODO(avidullu):** create it and paste the link.
209     - **TODO(avidullu):** add the escalation channel (email / WhatsApp group) for edge-case
    questions.
210     - NDA before any footage is shared ? template `NDA_TEMPLATE.md`. **TODO(avidullu):**
    have counsel
211         review, fill the blanks, and have Roora sign.
212     - Roora contact on file: info@rooragrp.com (HSR Layout, Bangalore).
213
214     > **? ADDED (fusion) ? data-handling guardrails (put in the brief / contract):**
215     > - **Exclude minors entirely.** If footage contains children, that video does **not**
    enter the
216     > labeled pool ? do not annotate it; flag it back to us.
217     > - **No biometric / identity labeling.** Refer to players only by **per-video
    pseudonyms** (e.g.
218     > "near side" / "far side") ? never a real name, and never a cross-video person ID. We
    do **not**
219     > collect gait or biometric annotations.
220     > - **No pose / skeleton / keypoint annotation** is requested.
221     > - **Work-for-hire IP assignment** (atop the NDA): the labels Roora produces are
    assigned to us as
222     > owned work product.
223
224     ## 9. Pricing request (please quote)
225
226     1. Price per rally and per minute of footage.
227     2. Pilot price (3 short videos, ~30?36 rallies) as a paid calibration round.
228     3. Indicative scale price for ~36?45 longer videos.
229     4. Turnaround for the pilot and throughput (rallies/day) at scale.
230     5. QA process and rework policy.
231     6. Whether annotators have racket-sport familiarity (helpful, not required).
232     7. Confirmation you can deliver the CSV schema above **(incl. the $A intake-manifest
    calibration columns).**
233
234     ---
235
236     # PART B ? Illustrated Annotation Guide (rater how-to)
237
238     **Sports:** Badminton, Pickleball, Table Tennis (singles AND doubles).
239
240     > **Illustrations:** boxes marked `[SCREENSHOT: ...]` are placeholders.
241     > **TODO(avidullu):** replace every `[SCREENSHOT: ...]` placeholder with a real frame
    grab from the
242     > pilot videos before sending this guide to Roora.
243
244     ## Part 1 ? What you are producing
245
246     One CSV per video (from the template Sheet), one row per rally:
247     `rally_number, start_mmss, end_mmss, start_time, end_time, shots_count, ending_reason,
    last_shot_outcome, score_before, score_after, notes`.
248
249     You type: `rally_number`, `start_mmss`, `end_mmss`, `shots_count`, `ending_reason`,
    `last_shot_outcome`
250     (+ optional score/notes). The `start_time` / `end_time` decimal-second columns FILL
    THEMSELVES via a
251     formula already in the template.
252
253     For every rally you record:
254     - When it started (serve contact) ? `start_mmss` (`m:ss.mmm`, e.g. `1:12.40`)
255     - When it ended (shuttle/ball dead) ? `end_mmss`
256     - How many shots happened ? `shots_count`
257     - Why it ended ? `ending_reason` · Where the last shot landed ? `last_shot_outcome`

```

```

258
259     Everything else (gaps between rallies) you IGNORE.
260
261     ## Part 2 ? The three golden rules
262
263     **RULE 1: IGNORE ALL DEAD TIME.** A rally is CONTINUOUS active play. The moment play
264 stops, the rally
265 is over. The time until the next serve is DEAD TIME and is never labeled.
266
267     DEAD          RALLY 1          DEAD          RALLY 2          DEAD
268 (ignore) |=====| (ignore) |=====| (ignore)
269           ^           ^           ^           ^
270         serve      shuttle      serve      shuttle
271        contact      lands      contact      lands
272       =start_1     =end_1     =start_2     =end_2
273     ...
274
275     Dead time includes: walking back to serve, picking up the shuttle/ball, wiping sweat,
276 bouncing the
277 ball before serving, the score being announced, replays/slow-mo, crowd or coach shots,
278 changeovers,
279 towel/medical timeouts, and PRE-MATCH / BETWEEN-GAMES WARM-UP (knock-up) ? see the Rule
280 2 caveat below.
281
282     **RULE 2: INCLUDE EVERY RALLY ? EVEN THE UGLY ONES.** Every served point gets a row,
283 including:
284 - Service faults (serve into net/out/illegal) ? usually `shots_count = 1`,
285 `ending_reason = service_fault`.
286 - Very short rallies (serve + one error).
287 - Lets / replays ? `ending_reason = let`; the replayed point is its OWN next row.
288
289     A missed rally is the worst error in this project ? the model must learn that short and
290 failed rallies
291 are still rallies.
292
293     > ?? **BUT WARM-UP / KNOCK-UP IS NOT A RALLY** ? a common trap, because Rule 2 says
294 "include
295 > everything." Pre-match and between-games warm-up LOOKS like rallying but is not scored
296 play; do not
297 > label it. How to tell it apart:
298 > - **No real serve.** Real play starts with a proper serve from the service position
299 (diagonal,
300 > behind the service line). Warm-up has players feeding cooperatively from mid-court
301 with no serve.
302 > - **Cooperative, not competitive.** Gentle hits straight to each other; nobody is
303 trying to win the
304 > point (no smashes-to-win, no scrambling/diving).
305 > - **Score isn't progressing.** No umpire / scoreboard not running / players casually
306 retrieve and
307 > re-feed; the score doesn't change. Often several shuttles/balls in use at once.
308 > - **Per sport:** badminton & pickleball ? real play starts with the underhand serve
309 from the service
310 > court / behind the baseline; table tennis ? real play starts with the tossed, two-
311 bounce serve.
312 >
313 > When unsure: find the first proper serve where the score starts to progress ? labeling
314 begins there.
315 > If you genuinely can't tell whether the match has started, flag the clip to us ? don't
316 guess.
317
318     **RULE 3: A SHOT = ONE CONTACT WITH THE RACKET / PADDLE / BAT.** Count every time the
319 shuttle/ball
320 touches a racket, paddle, or bat. The serve is shot 1.
321 - Count: any racket/paddle/bat contact (serve, return, smash, block?).

```

```

305     - Don't count: floor bounces, table bounces, net clips, the toss before a serve.
306
307     5-shot example: serve(1) ? return(2) ? hit(3) ? hit(4) ? winner(5) ? `shots_count = 5`,
308     `ending_reason = winner`.
309
310     ## Part 3 ? Where exactly does a rally start and end?
311
312     **START (`start_mmss`):** the exact frame the server's racket/paddle/bat CONTACTS the
313     shuttle/ball.
314     **Not the toss, not the wind-up ? the contact frame.** `[SCREENSHOT: serve-contact
315     frame]`
316
317     > **? NOTE (fusion):** this is the field that fixes the "held-shuttle" false positive ?
318     a player
319     > holding/flicking the shuttle before serving is **dead time**, so it never gets a row.
320     Faithful
321     > serve-contact boundaries are exactly what the detector is graded against.
322
323     **END (`end_mmss`):** the exact frame the shuttle/ball becomes DEAD:
324     - Badminton: shuttle touches the floor, or play stops on a fault.
325     - Pickleball: the second floor bounce, or the frame the ball hits the net / lands out.
326     - Table tennis: the second bounce on one side, a non-serve net hit, or off the table.
327
328     `[SCREENSHOT: end frame]`
329
330     If occluded/off-screen, pick the closest clear frame and add a note. **Boundary
331     tolerance:  $\pm 0.3$  s
332     (~ $\pm 9$  frames at 30 fps).**
333
334     ## Part 4 ? `ending_reason`: how to choose
335
336     Pick exactly ONE value for how the rally ended. Stop at the first match:
337     1. Did the SERVE itself fail / was illegal (and ended the point)? ? `service_fault`
338     2. Was the point REPLAYED (let / interruption)? ? `let`
339     3. Did the final shot land IN and the opponent make NO legal contact? ? `winner` (a
340     clean winner)
341     4. Did the losing player make an error (netted / out / missed) UNDER PRESSURE ?
342     stretched,
343     wrong-footed, or rushed by a strong shot? ? `forced_error`
344     5. Did the losing player make an error on a ROUTINE, UNPRESSURED ball they had the time
345     and position
346     to make? ? `unforced_error`
347     6. None of the above / cannot tell ? `other` (add a note)
348
349     **FORCED vs UNFORCED ? the one judgment call.** Decide using the opponent's PRECEDING
350     shot:
351     - `forced_error` = the opponent's shot CAUSED the miss (pace, placement, deception; the
352     player barely
353     reached it or was pulled out of position).
354     - `unforced_error` = a clean, makeable ball the player missed on their own.
355     - When genuinely unsure between the two, choose `unforced_error` and add a brief note.
356
357     Notes:
358     - A clean winner means NO legal contact by the opponent. If they got a racket/paddle/bat
359     on it but
360     netted or hit out, that is a `forced_error` / `unforced_error` (their shot), not a
361     winner.
362     - A serve into the net is `service_fault` (rule 1 beats everything).
363
364     **THEN SET `last_shot_outcome`** ? where the last shot physically went. This is a
365     SEPARATE,
366     observational column (no judgment): just look at the last shot.
367     - `in` ? landed in (winners; also a put-away the opponent couldn't reach).
368     - `net` ? went into the net.
369     - `out` ? landed outside the court / past the lines.

```

```

357     - `n_a` ? not applicable: use for `service_fault` and `let`.
358
359     So each rally gets both: WHY (`ending_reason`) and WHERE (`last_shot_outcome`) ? e.g. a
smash winner =
360     `winner, in`; a clear forced long = `forced_error, out`; a routine ball netted =
`unforced_error, net`.
361
362     ## Part 5 ? Per-sport detail
363
364     ### Badminton
365     - **Serve:** underhand, below the waist, diagonal. Start = racket?shuttle contact.
366     - **Rally end:** shuttle hits the floor, or play stops on a fault.
367     - **Shots:** every racket?shuttle contact; serve = shot 1.
368     - **Service faults** (1-shot rally, `service_fault`): serve into net, serve outside
service court,
369     contact above waist, foot fault.
370
371     Worked example: `[SCREENSHOT: badminton serve]`
372     - Rally 4: serve `1:28.30`, long rally, smash winner, shuttle lands `1:44.10`, 12
contacts.
373     `4, 1:28.30, 1:44.10, (auto), (auto), 12, winner, in, 3-5, 4-5,`
374     - Rally 5: serve `1:57.00` straight into the net.
375     `5, 1:57.00, 1:57.90, (auto), (auto), 1, service_fault, n_a, 4-5, 4-6,`
376
377     **Badminton ? doubles (2 per side).** All rally rules are IDENTICAL to singles:
378     - start = serve contact; end = shuttle down / fault; a SHOT = any racket?shuttle contact
by EITHER
379     player on a side.
380     - `shots_count` = TOTAL contacts across both teams in the rally. Do not track which
player hit ? just
381     count every contact.
382     - Net exchanges (drives, flat pushes) come fast ? frame-step so you neither miss nor
double-count.
383     - Only the designated receiver returns the serve; if the wrong player returns, the rally
ends on a
384     fault ? still a row (`service_fault` if it's a serve/receiving-position fault, else
`other` + note).
385     - Two teammates may NOT both hit the shuttle in succession (double-hit). Two contacts on
the SAME side
386     back-to-back = a fault ending the rally ? count both contacts, then the rally ends.
387
388     ### Pickleball
389     - **Serve:** underhand, paddle below waist, hit diagonally; must clear the non-volley
zone (the
390     "kitchen"). Start = paddle?ball contact.
391     - **Two-bounce rule:** the serve must bounce once in the receiver's court and the return
must bounce
392     too, before either side may volley. These bounces are NORMAL play, not rally ends, and
bounces are
393     NOT shots.
394     - **Rally end:** ball bounces twice, lands out, hits the net (fails to cross), or a
fault (e.g.
395     volleying in the kitchen).
396     - **Shots:** every paddle?ball contact only; serve = shot 1; floor bounces are not
shots.
397     - **Service faults** (`service_fault`): serve into net, out/long, short in the kitchen,
foot fault.
398
399     Worked example: `[SCREENSHOT: pickleball serve + kitchen line]`
400     - Rally 2: serve `0:33.10`; serve bounces, return bounces, a few volleys; receiver nets
an easy one at
401     `0:41.85`; 6 paddle contacts.
402     `2, 0:33.10, 0:41.85, (auto), (auto), 6, unforced_error, net, 1-0, 2-0, easy ball
netted`
403

```

```

404    **Pickleball ? doubles (the standard format, 2 per side).** All rally rules IDENTICAL to
singles:
405    - start = serve contact; a SHOT = any paddle?ball contact by EITHER teammate; floor
bounces still not
406    shots; two-bounce rule still applies.
407    - `shots_count` = TOTAL paddle contacts across both teams; don't track which player.
408    - Serving sequence (FYI only ? you do NOT annotate it): both partners serve before the
serve passes to
409    the opponents, except the very first service turn of the game. The called score has
three numbers.
410    Ignore for annotation; just mark boundaries / shots / ending.
411    - If both teammates contact the ball in succession, that's a fault ? rally ends; count
the contacts,
412    set the ending per the fault.
413
414    ### Table tennis
415    - **Serve:** ball tossed up from an open palm, then struck so it bounces once on each
side. Start = the
416    racket?ball strike (not the toss).
417    - **Rally end:** a player fails a legal return ? ball double-bounces on their side, goes
off the table,
418    or is hit into the net on a non-serve shot.
419    - **Shots:** every racket?ball contact only; serve = shot 1; table bounces are not
shots.
420    - **Let:** a serve that clips the net but is otherwise legal ? the point is replayed.
421    `ending_reason` = let`, and the replay is its own next row.
422    - **Service faults** (`service_fault`): missed/illegal toss, serve into net, serve off
the table.
423    - **Speed warning:** contacts can be < 0.2 s apart ? step frame-by-frame to count shots.
424    - **Doubles** (if a clip appears): partners must hit in STRICT ALTERNATION and serve
diagonally; a
425    missed alternation is a fault that ends the rally. Same rule ? count every racket?ball
contact as a
426    shot; `shots_count` is the team total.
427
428    Worked example: `[SCREENSHOT: table-tennis serve contact]`
429    - Rally 9: serve `4:10.40`; fast 7-contact rally; forehand winner, opponent no contact;
dead `4:14.05`.
430    `9, 4:10.40, 4:14.05, (auto), (auto), 7, winner, in, 6-9, 6-10,`
431    - Rally 10: serve `4:22.10` clips the net (otherwise good) ? let.
432    `10, 4:22.10, 4:23.00, (auto), (auto), 1, let, n_a, 6-10, 6-10, net-cord serve,
replayed`
433
434    ## Part 6 ? Edge cases & FAQ
435
436    - **Two rallies with almost no gap?** Still two rows. `end_time[N] <= start_time[N+1]`,
never overlap.
437    - **Interrupted & replayed?** Label the interrupted attempt as its own row
(`ending_reason=let`), then
438    the replayed point as the next row.
439    - **Serve tossed but caught / re-served (no contact)?** No contact = no rally yet. Start
at the actual
440    serve contact; don't label the aborted toss.
441    - **Opponent reached the ball but netted/hit out?** That's a
`forced_error`/`unforced_error` (their
442    shot), not winner ? forced if your shot pressured them, unforced if it was routine.
443    - **Can't read the scoreboard?** Leave `score_before`/`score_after` blank. Never guess.
444    - **"Let" where the rules say play on?** If play continued, it's NOT a let ? keep one
rally and judge
445    the real ending.
446    - **Warm-up / knock-up / practice points?** Not real rallies ? do not label them. Start
at the first
447    real, scored serve. See the Rule 2 caveat (Part 2). When unsure whether the match has
started, flag
448    the clip ? don't guess.

```

```

449 - **Doubles?** Same rules ? a shot is any one racket/paddle contact by whichever player
hits it. Never
450     attribute `shots_count` to an individual, just total the contacts.
451 - **Server completely whiffs (no contact)?** Rare. True zero contact ? `shots_count` =
0`,
452     `ending_reason` = service_fault`, add a note. A graze = 1 contact = 1 shot.
453
454 ## Part 7 ? Self-check before you submit (per video)
455
456 - [ ] `rally_number` is 1,2,3? no gaps or repeats.
457 - [ ] Every `start_time` < `end_time` (check the auto-filled decimal columns populated).
458 - [ ] No overlaps: `end_time[N]` <= `start_time[N+1]`.
459 - [ ] Times entered as `m:ss.mmm`; decimal `start_time`/`end_time` auto-filled and look
right.
460 - [ ] No dead-time rows (scan the gaps ? dead time, not missing rallies).
461 - [ ] Every served point present, including service faults and lets.
462 - [ ] `shots_count` looks right (re-step any rally you rushed; doubles = team total).
463 - [ ] `ending_reason` in {winner, forced_error, unforced_error, service_fault, let,
other}; notes
464     filled when `other`.
465 - [ ] `last_shot_outcome` in {in, net, out, n_a} (`n_a` for service_fault/let).
466 - [ ] File downloaded as `.csv`, header row present, UTF-8.
467 - [ ] Did NOT trim / re-encode / rename the source video.
468 - [ ] **? (fusion) ** intake-manifest row filled: `court_corners`, `net_line`,
`singles_or_doubles`.
469
470 ## Part 8 ? Mini glossary
471
472 - **Rally** ? one continuous passage of play from serve contact to the shuttle/ball
dead.
473 - **Shot** ? one contact between the shuttle/ball and a racket/paddle/bat (either
teammate in doubles).
474 - **Dead time** ? any time the shuttle/ball is not in active play. Never labeled.
475 - **Service fault** ? a failed/illegal serve that ends the point. Still a labeled rally.
476 - **Let** ? a replayed point. Labeled as its own rally with `ending_reason=let`.
477 - **Winner / forced_error / unforced_error** ? the final shot was a clean winner
(opponent made no
478     legal contact) / the loser's error was forced by a pressuring shot / the loser's error
was on a
479     routine, unpressured ball.
480 - **start_time / end_time** ? decimal seconds from the start of the provided file (auto-
filled from mm:ss).
481 - **? court_corners / net_line / singles_or_doubles ** ? per-video calibration in the
intake manifest (§A).
482
483 ---
484
485 ## Appendix B ? Tier-B extensions (NOT in the day-1 deliverable) ? ADDED (fusion)
486
487 > Pilot these on the 3-video calibration round, then renegotiate price before scaling.
**Do not start
488 > these without explicit instruction ** ? they are listed so Roora can quote them.
489 >
490 > - **Per-shot hit timestamps** ? a timestamp for *each* racket/paddle/bat contact (not
just the
491 >     `shots_count` total). Extends the §7 CVAT method (box every contact, not only the
two boundary
492 >     frames). Trains/validates the audio hit-detector and shot-cadence features.
493 > - **Per-shot type** ? for each contact, a label from {serve, clear, drop, smash, net,
drive, lift}.
494 >     Needs its own short guideline + an inter-annotator agreement check.
495

```

Now I'll read the archives/quality-iterations directory to see what's there:

18. user (2026-06-16T03:12:26.277Z)

```
1      # quality-iterations/
2
3      This directory holds completed phases of quality improvement work.
4
5      As we receive new labeled videos and iterate on model accuracy, recall/precision, audio
6      fusion, candidate quality, etc., each major iteration or batch of improvements will be captured
7      here once it reaches a terminal state (DONE / COMPLETED / REJECTED).
8
9      See `../README.md` (under archives/) for the overall archival policy.
10
11     ## Organization
12     - Dated subdirectories, e.g. `2026-06-initial/`
13     - Each iteration should include:
14       - A README.md with STATUS, trigger (new data, specific goals), key changes, and links
15       to relevant checkpoints under `../checkpoints/`.
16       - References to updated eval baselines, regression results, report enhancements, etc.
17
18     This keeps the live `docs/` clean while preserving a rich historical trail for
19     regression detection across quality work.
20
21     ## Iterations
22     - **`2026-06-initial/`** ? SEED. Opened the quality phase ahead of the awaited labeled-
23     video batch (status: in progress; superseded for first results by the iteration below).
24     - **`2026-06-human-gt-tuning/`** ? **COMPLETED (2026-06-08).** First eval against
25     **human** ground truth on a partially-labeled video: built the partial-label eval harness +
26     labeled-window methodology, produced the first trustworthy F1 (baseline 0.650 ? tuned 0.742,
27     *fit-on-self*), per-rally failure analysis, a config sweep, and ? via adversarial review ? found
28     + fixed a real ingestion bug (negative-start import abort). Cross-repo: `sports-data-collector`
29     PR #13.
30     - **`2026-06-multivideo-lovo/`** ? **COMPLETED (2026-06-08 ? extended to n=6,
31     2026-06-09).** Built a **6-match** human corpus (3 singles incl. Mahadevpura, 1 doubles, 2
32     multi-court; 3 cameras/frame-rates) + cross-video **leave-one-video-out**. Key finding: the
33     unsupervised windowing heuristic is **footage-fragile** ? it fails on **multi-court** footage
34     (background games; Kushagra 0.16 / gBaaddy 0.28) where the shuttle signal is present (97% of
35     rallies have motion) but rally-vs-dead-time isn't velocity-separable. **Prototyped a learned
36     per-frame segmenter** (`backend/eval/rally_seq_proto.py`): held-out AUC **0.83?0.92 on every
37     match, decoupled from the multi-court confound**. **The learned model OVERTOOK the heuristic at
38     n=3 and holds ~+0.10**: n=6 LOVO mean **learned 0.583 vs heuristic 0.480**. The **"contrast"
39     metric** (in-rally÷dead-time shuttle visibility) predicts heuristic F1 monotonically across 5
40     matches. Data-backed green light for `../research/OWNED_MODEL_TRAINING_STUDY.md`.
41     - **`2026-06-first-ab-experiment/`** ? **COMPLETED (2026-06-14).** The **first end-to-
42     end A/B** on the n=6 golden corpus through the experiment harness + the C1?C4 measurement-
43     honesty helpers, run **100% offline** (re-scored persisted WASB trajectories, no GPU). One
44     consistent candidate set re-scored three ways: **Gate-A (net-crossings ? 2) is a real win** ?
45     ?F1 +0.074, **6/6 matches, sign-test p = 0.031, CI [+0.042, +0.104]**, over-segmentation
46     (segment-count ratio) **1.68 ? 1.01**; the **window-level 5-feature logistic does NOT clear the
47     bar** (?0.081, 2W/4L, a threshold/calibration artifact ? distinct from the stronger per-frame
48     prototype above). The headline is methodological: **C1 confidently promoted one variant and
49     refused another**, and **C4 surfaced over-segmentation tIoU hid** ? the framework held up on
50     first real use.
```

19. assistant (2026-06-16T03:13:38.777Z)

Perfect. Now let me compile the comprehensive survey. I have enough information to provide a detailed breakdown.

SURVEY: Rally-Detection Design Docs ? Status, Key Decisions, Measured Results

Document Index & Status

```
| Document | Status | Coverage | Key Decisions |
|---|---|---|---|
| **README.md** | Living index | Documentation roadmap, cross-references | Indexes all active
docs; hardening loop marked COMPLETED (2026-06-14) |
| **CODE_MAP.md** | Living reference | Repo structure, contracts, file placement | Canonical
source for file organization; 2026-06-10 audit sweep added ~8 behavioral changes + 8 new tests |
| **EXPERIMENT_HARNESS.md** | **P1 SHIPPED** (2026-06-13) | A/B + shadow eval, LOVO-honest
measurement, promotion gate | `Variant` abstraction, `compare()`, `compare_unlabeled()`,
experiment leaderboard; **zero-regression by design** (cues default OFF, promote only on
measured lift) |
| **MULTI_SIGNAL_FUSION_PLAN.md** | **P1 DONE** (2026-06-13) | Shuttle + person + audio, one
fusion layer | Person cue extracted to reusable `PersonDetector`; P2 `FusionSegmenter` shipped
(default-OFF); **licensing all MIT/Apache/BSD** |
| **ANALYZERS_AND_RECONCILERS.md** | **Design ? owner-gated, NOT started** | Signal-fusion
framework; frame/window/session scopes; no special-casing spine | Multi-agent literature review
($13) completed; identifies standard vs. novel patterns; course corrections C1?C13 documented |
| **HOW_RALLY_DETECTION_WORKS.md** | Living conceptual map | 2-layer model (detection vs.
segmentation) | Clarifies shuttle vs. person cues; open-question list includes fusion, per-video
calibration, learned segmenter |
| **QUALITY_ITERATIONS.md** | **Living log** | Reverse-chronological quality moves, measured
effects | Baseline table (heuristic 0.480 ? Gen-0 0.626 ? Gemini 0.93); lessons carry-forward
(small-N discipline, LOVO, learned model needed) |
| **OWNED_MODEL_IMPLEMENTATION_PLAN.md** | **Phase-0 PLAN** | Roadmap to owned-model +
multisport | Framework = ms-swift (Apache-2.0); base = Qwen3-VL (Apache-2.0); **START WITH TAL
HEAD, NOT VLM** | licensing constraints (MIT/Apache/BSD only; no paid-LLM training) |
| **NEXT_STEPS.md** | **Living** | Prioritized next moves, pointers | M0 in-container quick wins
COMPLETE; GPU-box P3 measurement next; corpus growth = single highest-leverage |
| **ROORA_LABELING_GUIDELINES.md** | **v8 (2026-06-13)** | Vendor annotation spec + rater guide
| Court calibration, singles/doubles, intake manifest; 3 non-negotiables (ignore dead time,
include every rally, one shot = one contact) |
| **PERSONALIZATION_PLAN.md** | **Design ? HYBRID adopted** | Per-user tuning on generalization
baseline | Owner decisions locked (2026-06-14): identity (design north-star, implement local),
`min_n` = 5 (promotion floor, not service gate), Gate-A-only structural guard |

---
```

CRUCIAL: \$13 Literature Review (ANALYZERS_AND_RECONCILERS.md)

****What the multi-agent literature review surveyed and the verdicts:****

Patterns Surveyed (\$13.1, Verdict Table)

```
| Pattern (our framing) | Field's term | Closest prior art | Verdict |
|---|---|---|---|
| Candidate-windows + scorer + reconciler | Temporal Action Localization/Segmentation (TAL/TAS)
+ boundary refinement | ASRF (arXiv:2007.06866), BMN, Activity Grammars | **partial-overlap** |
| "Signal not verdict" ? learned arbiter | Late/score-level fusion + stacked generalization |
Snoek 2005, Wolpert 1992, CVIU 2009 racket-sports | **partial-overlap** |
| Detection?decision, persist-once/re-score | Frozen-backbone cached features + decoupled
proposal/classification | S-CNN, CBR, Decouple-SSAD | **partial-overlap** (? field default) |
| **Report-first reconciler** (contradiction ? audited correction, measured before apply) |
Model/consistency assertions, constraint-based repair | Model Assertions (2003.01668),
HoloClean, TFDV, EventAnchor | **novel-combination** |
| Scarce interval labels + tiny LR + LOVO-by-match | Point-supervised TAL + grouped (LOGO/LOSO)
CV | SF-Net, LOOCV-bias (arXiv:2406.01652), Snorkel | **partial-overlap** |
| Multi-scope (frame/window/session); warm-up conditions per-window | Multi-scale/hierarchical
modeling + play-break + phase recognition | Tjondronegoro 2003, GL-MSTCN, MS-TCN | **partial-
overlap** |
```

Verdict Summary (\$13.1 TL;DR)

> **"The architecture's skeleton is standard and must be *cited*, not *claimed*: candidate-
windows + learned scorer = *Temporal Action Localization* (proposal-then-score); reconciler**

over-segmentation fixes = **Temporal Action Segmentation** refinement; "signal not verdict / learned arbiter over cue columns" = **late / score-level fusion** realized as **stacked generalization** ? ~20 years old and already present in racket sports. The defensible delta is the engineering/governance envelope: detection?decision persisted as a re-score substrate + a **report-first, A/B-lift-gated, idempotent, human-in-the-loop reconciler-as-linter** keyed on decision-vs-evidence contradiction."

Only the reconciler-as-linter and warm-up-as-soft-session-feature rated `novel-combination`; the rest are `partial-overlap`.

In-Domain Prior Art Already Hard-Codes These Cues (Our Delta = No-Branch Discipline)

- Non-broadcast badminton rally start/end (See et al. 2024/2025)
- Shuttle-state/hit detection (MonoTrack arXiv:2204.01899; Hsu/Yu/Cheng, Sensors 2024)
- Service-fault (Goh et al., Sensors 2023)
- Tennis in-play state machines
- Badminton point-segmentation (Ghosh et al. 2018)
- ShuttleSet (arXiv:2306.04948)

The delta is not cues ? it's the *no-branch, weight-learned-cues* discipline + the measured-before-apply reconciler.

Prior Measured Results

Quality Baseline Table (QUALITY_ITERATIONS.md, NEXT_STEPS.md)

Path	Mahadevpura (clean)	Kushagra (multi-court)	Notes
Heuristic (velocity windowing)	0.657	0.157	**footage-fragile**; dead-time confusion on multi-court
Gen-0 owned head (LOVO, n=6)	0.744	0.561	LogReg on 5 trajectory features; threshold-in-LOVO needed (C2)
Two-stage WASB+Gemini refine	0.83	?	Cost-efficient for long tapes
Gemini wrapper (gemini-2.5-flash)	**0.93**	**0.66**	**Commoditized path for clean footage; multi-court target for owned model**

Reading: Own model's ship target = **0.66 on multi-court** (the contrast-confound problem); Gen-0 is 0.10 behind on multi-court, **corpus-growth-sized gap, not architecture-sized**.

Gate-A Measurement (QUALITY_ITERATIONS.md, "Acting on the First A/B")

- **?F1 +0.074** (6/6 golden matches, sign-test **p=0.031**, CI [+0.042, +0.104])
- **Over-segmentation: segment-count ratio 1.68 ? 1.01**
- **Real-rally windows lost: 12 windows across 6 videos on served COARSE windowing** (net_crossings ~1.0; the gate's threshold of ?2 is too strict for the served path)
- **Eval ? serve lesson:** the +0.074 is conditional on the eval harness's recall-oriented *proposal* windowing (0.005 velocity, 1.0 merge-gap, 0.5 min-duration); served windowing is coarse (0.02, 2.0, 2.0) and already merges aggressively. **Gate-A flipped to default and then REVERTED** on served-side regression check (**served ?F1 ?0.008, ~12 real rallies dropped**). Now **opt-in** (`indexing.fusion.min\_crossings: 2`), default OFF.

Person + Audio Fusion Measurement (QUALITY_ITERATIONS.md, "Fusion P3 + P4 MEASURED")

Variant (LOVO, n=6, IoU ? 0.5)	F1	?F1	95% CI	Sign test	CI clears 0	Status
Shuttle-only (CONTROL)	0.307	?	?	?	?	?
+ person (P3)	0.286	**?0.021**	[?0.165, +0.162]	2W/4L (p=0.688)	**No**	Default-OFF
+ person + audio (P4)	0.366	**+0.079**	[?0.054, +0.203]	4W/2L (p=0.688)	**No**	Default-OFF (promising, not promotable)

Person adds noise (no court polygons ? `players\_in\_court` counts spectators); **audio is the more promising lead** (+0.079 / +27.7%, wins 4/6 videos), but **n=6 can't confirm it**. Both

stay default-OFF (promote-only-on-lift discipline). Recorded in
`eval\_baselines/experiment\_leaderboard.json`.

Personalization Phase-0 (QUALITY_ITERATIONS.md checkpoint entry, 2026-06-14)

Three **honest-eval RAILS** that keep per-user tuning from promoting noise:

1. **Per-user promotion gate** (PR #141): min-N floor (5 videos) + structural-guard exemption (Gate-A stays on)
2. **Held-out-USER learning curve** (PR #142): falsifiable "<K videos to superior quality" metric
3. **Per-user no-regression guard** (PR #143): incumbent CONTROL; candidate never regresses user

Stated Eval Approach & Key Lessons

Evaluation Framework (EXPERIMENT_HARNESS.md §3; EVALUATION.md; QUALITY_ITERATIONS.md)

- **LOVO** (Leave-One-Video-Out): score on held-out *video*, never held-out *window* (windows in one video are correlated). Group by **match** (group-aware LOGO/LOSO).
- **IoU threshold**: default 0.5 temporal IoU for matching + scoring.
- **Honest stats** (C1?C4, QUALITY_ITERATIONS.md):
 - **C1**: Bootstrap 95% CIs + paired sign test (never a single mean at n?6); RLOOCV for bias-aware folds
 - **C2**: Calibrate each producer's confidence (Platt/isotonic) before the scorer
 - **C3**: Out-of-fold stacking (learned producer ? leave-one-group-out outputs to the LR)
 - **C4**: Report field metrics (F1@{10,25,50}, edit score, tight/loose mAP) alongside temporal-IoU; tIoU alone is blind to over-segmentation
- **Eval ? serve lesson**: the experiment harness windowing (recall-oriented, permissive candidates) differs from the served path (coarse, few candidates). A harness win doesn't predict served behavior ? align windowing or re-run promotions on served windowing.

"Signals?Learned-Scorer, NO Special-Casing" Spine (ANALYZERS_AND_RECONCILERS.md §1)

- > "Every producer emits a SIGNAL carrying a value + confidence. The SCORING MODEL consumes those signals and learns to weight them. We do NOT add hand-coded `if/else` branches into the decision."
- The **SVM/LogReg learned scorer** (ADR-004, `backend/eval/classifier.py`) is the only place a decision is made.
 - Analyzers propose signals as feature columns; the scorer learns weights.
 - Reconcilers run *outside* the decision path as a post-hoc linter (report-first, never silent).
 - **Invariant**: all analyzer flags OFF + reconciler report-only ? candidate rows AND predictions byte-identical to control.

No-Regression Guarantees (EXPERIMENT_HARNESS.md §6)

1. **New cues default OFF** ? per-cue enable flag; merged code can't change behavior until explicitly enabled.
2. **Control variant pinned** ? frozen recipe + classifier model hash, always registered & default.
3. **Promotion only via the eval gate** ? `run\_regression.py` in CI (exit 1 blocks); no silent default change.
4. **Experiments are records** ? variant-spec hash ? per-video + LOVO scores + label-set hash + timestamp, persisted to JSON leaderboard.
5. **Decoupled events** ? "merge an experiment" and "change the default" are separate, independently revertible actions.

Owned-Model Plan: Base Model & Licensing Constraints

Design Principle (OWNED_MODEL_IMPLEMENTATION_PLAN.md §0)

> **"Envision now, build later. Every deferred piece must be architecturally envisioned now so we never hit a mammoth refactor. Design the seams (data contracts, export boundaries, event store) up front and implement behind them when each milestone is greenlit."**

Hard Guardrails (Non-Negotiable)

1. ****Commercial-clean licensing for EVERY dependency ? code, data, AND weights: MIT / Apache-2.0 / BSD only.**** Reject viral/NC/custom-restrictive (Gemma 3, Llama 4 700M-MAU cap, Commons-Clause).
2. ****Never train on paid-LLM (Gemini/OpenAI/Anthropic) outputs.**** Also forbidden: public grounding-CoT datasets that were Gemini-generated (e.g. TVG-R1's 56K).
3. ****Exclude minors; per-user training data needs explicit opt-in; split data by match.****
4. ****? Base-model pretraining provenance is unauditible for *every* open VLM ? explicit owner risk-acceptance required.**

Model Strategy (Verified 2026-06-09)

Component	Choice	Rationale	
--- --- ---			
Framework	**ms-swift (Apache-2.0)**	One forge for video + audio + timestamp-grounding + GRPO/RFT; keep **unsloth** (Apache-2.0) for 8 GB local SFT/compression (on-device goal)	
Base (VLM)	**Qwen3-VL (Apache-2.0)**	Long-video context, timestamp emission, multi-turn chat; **InternVL3** (MIT/Apache) = fallback	
Gemma-4	AMBER (bench-only)	Apache grant likely but Prohibited-Use unconfirmed; video/timestamp capability disputed	
START POINT	**TAL head, NOT VLM**	Fine-tuned VLMs ?50% R@0.7 on open-domain; our head = de-risked AUC 0.83?0.92, trains in minutes, doubles as pseudo-label teacher + honest baseline VLM must beat; **head now ? VLM later**	

Conversational QA Architecture (Seam Built Now, Model Deferred)

****Punt the conversational *feature* until extraction quality ? "very high." But build the *seam* now:****

- ****Extract once, query many:**** TAL head ? later VLM ? per-video ****EVENT STORE**** ? ****structured QUERY layer**** (text-to-SQL + ffmpeg clip-cut)
- ****Per-video event-index schema (first-class data contract):**** `{video_id(MD5), rally_ordinal, sport, start, end, shot_count, ending_reason/fault_tags, derived_clip_ref, provenance, confidence}`
- ****Reserve fault-tag taxonomy now**** (e.g. ``service_fault``), even if unpopulated
- ****Paid Gemini may answer open-ended edge questions at inference**** (allowed ? it reasons over our data, not trained on its outputs)

Multisport (Single Sport-Agnostic Model)

- ****Model/framework: UNCHANGED**** ? one shared TAL head + Qwen3-VL, fine-tuned across sports
- ****What multisport changes = DATA + DETECTORS (scale up):**** each sport needs labeled matches (split by match) + clean trajectory detector
- ****Binding blocker:**** clean MIT/Apache/BSD shuttle/ball detectors for tennis/TT/pickleball/padel not yet identified; WASB has unresolved C3 checkpoint-provenance issue
- ****? Single-model-multisport-temporal-generalization is OPEN HYPOTHESIS****, not proven (badminton-only corpus; RacketVision refuted on verification)

Phase Structure & Status

Core Quality Track (Default-OFF, Promote-Only-On-Measured-Lift)

****Shipped (PR #112 + 2026-06-13):****

- P1: Person cue extracted to reusable ``PersonDetector`` (PR #115)
- P1: Experiment harness (``compare()``, ``compare_unlabeled()``, leaderboard) ? shipped 2026-06-13 ?
- P2: ``FusionSegmenter`` (Gate A signal + person-feature persistence), default-OFF ? shipped

2026-06-14 ?

****In progress / owner-gated:****

- P3: Feature-level (learned) fusion ? person features + LOVO lift measurement
- P4: Audio cue integration (hit_detector already scaffolded in ADR-007)

Hardening Loop (COMPLETED 2026-06-14)

- ****T1?T5 positive tests**** for 5 identified issues + multi-layer regression protocol
- ****626-test suite**** + ablation/experiment layer + pre/post-change verification checklist
- ****Archived:**** `docs/archives/CODE\_AUDIT\_AND\_TEST\_HARDENING-COMPLETED-2026-06-14.md` +
`docs/archives/HARDENING\_LOOP\_HANDOFF-COMPLETED-2026-06-14.md`

Critical Cross-Links for New Foundational Research

To avoid duplication, a new foundational doc should:

1. ****Cite §13 literature review explicitly**** ? don't re-derive the skeleton (TAL/TAS/late-fusion) or in-domain prior art
2. ****Anchor on the "no-branch spine"**** (ANALYZERS_AND_RECONCILERS.md §1) ? the core governance principle
3. ****Reference the eval ? serve lesson**** (QUALITY_ITERATIONS.md, Gate-A revert) ? the honest pitfall
4. ****Build on the framework's three temporal scopes**** (frame/window/session) ? the reusable architecture
5. ****Ground course corrections C1?C4 and Tier-2/3**** (§13.4?C13) ? the open research debts
6. ****Carry the small-N discipline**** (few features, logistic regression, LOVO-by-match, not held-out-window)
7. ****Inherit the licensed-only, no-paid-LLM-training, minors-excluded guardrails****
(OWNED_MODEL_IMPLEMENTATION_PLAN.md, ROORA_LABELING_GUIDELINES.md §8)

This survey is complete and all document/section anchors are canonical (file paths + §X references as they appear in the codebase).