

claude-windows / 6a82308f-69d6-43a8-9ea4-d8eaf497253d

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6a82308f-69d6-43a8-9ea4-d8eaf497253d\subagents\workflows\wf_3d536535-241\agent-ac08d9e2b188044fd.jsonl`  
- SHA-256: `4e755c1effb41ed1122363925466d500f05eafea9db1a39e85001734b9912e41`  
- Source modified: `2026-06-16T03:20:34+00:00`  
- Imported at: `2026-07-05T16:48:27+00:00`  
- project: `wf_3d536535-241`  
- session_id: `6a82308f-69d6-43a8-9ea4-d8eaf497253d`
```

Transcript

1. user (2026-06-16T03:19:04.326Z)

Source Extractor

Research question: "Research question: **What concrete, recent (2022?2025) techniques and evaluation practices should a local, commercially-licensed badminton rally-detector adopt to close the quality gap with a strong cloud VLM (Gemini), given a tiny labeled set (~6 ? ~16 golden videos)?** Produce an adversarially-verified, cited report organized around the four dimensions below. Prefer specific named methods, papers (arXiv/venue), and reported numbers over generic advice; flag anything you cannot verify.

CONTEXT (so you target the real gap, not generic ML):

- System: a local pipeline detects the shuttle per-frame (WASB / TrackNet, MIT-licensed) ? a heuristic turns the trajectory into candidate rally windows ? today an optional Gemini "handoff" confirms rallies. We just measured a fully-local (no-AI) run vs Gemini on 3 matches: local emits far FEWER but MUCH LONGER windows (mean 65s vs 6.1s; one window spanned a whole 174s clip) ? it does NOT miss activity, it FAILS TO SEGMENT. Root cause (code-confirmed): a frame is "active" iff the shuttle is visible AND moving above a small per-second-normalized velocity threshold; active frames within a 2 s gap are merged; there is NO max-window cap and NO inactivity/boundary split. So "shuttle is moving" is conflated with "rally in progress," and dead-time + serves + knock-ups bridge rallies into mega-windows. Shuttle tracking itself is excellent (96?99% per-frame visibility).

- HARD CONSTRAINTS: weights/data/code must be MIT/Apache-2.0/BSD only (no viral/NC/custom licenses, reject Gemma-3/Llama-4-MAU-cap); we may NEVER train owned weights on paid-LLM (Gemini/OpenAI/Anthropic) outputs (terms/copyright), though paid LLMs may be used at inference; minors excluded from any training pool; runs on a single local GPU.

- WHAT WE ALREADY HAVE (do NOT re-survey the basics ? go deeper/newer): the architecture skeleton is known to be standard Temporal Action Localization/Segmentation (proposal+scorer) + late/score-level fusion + stacked generalization; in-domain prior art (MonoTrack, ShuttleSet, See et al. 2024/25 non-broadcast badminton rally start/end, Ghosh et al. point-segmentation, tennis/TT play-break state machines) is already cited. We have a working eval harness: leave-one-video-out (grouped by match), temporal-IoU F1 @0.5, F1@{0.1,0.25,0.5}, mAP@tIoU, segment-count-ratio, bootstrap 95% CIs, paired exact sign-test, Platt/isotonic calibration + Brier/ECE, a tiny numpy LogisticRegression scorer over per-window feature columns, an experiment leaderboard, and an "eval?serve" guard (a cue measured +0.074 F1 on permissive proposal windowing but ?0.008 on the coarser served windowing ? reverted). Available but default-OFF per-window signals: net-crossing count (rally gate), 5 person/court features, 3 audio-hit features.

THE FOUR DIMENSIONS:

1) WINDOWING & BOUNDARIES (Temporal Action Localization/Segmentation): beyond proposal+score ? what are the best *recent* methods for (a) turning a dense per-frame signal into well-segmented actions, (b) explicitly controlling OVER-segmentation/merging, (c) boundary detection/regression and splitting merged segments? Cover e.g. MS-TCN/MS-TCN++, ASRF (action-boundary regression), BMM/BSN boundary-matching, ActionFormer/TriDet/TemporalMaxer (actionness + boundary heads), smoothing/boundary losses, and the standard metrics for over-segmentation (segmental F1@k,

edit/Levenshtein score, mAP@tIoU). Which transfer to a 1-D shuttle-trajectory + auxiliary-cue setting on tiny data?

2) SHUTTLE-TRAJECTORY ? RALLY SIGNAL PROCESSING (racket-sports specific): how do published systems convert ball/shuttle tracking into rally/point segments and distinguish "in-play" from dead-time? Cover rally-state machines, serve/let detection, net-crossing & two-sided-occupancy cues, hit/contact (audio + kinematic) detection, and fps-invariant trajectory features. What features best separate "rally" from "shuttle merely moving"?

3) SMALL-DATA STRATEGIES: with ~6?16 labeled videos, what gives the most quality per label ? active learning (which videos/segments to label next; uncertainty/diversity/core-set), semi-/weakly-/point-supervised TAL (e.g. SF-Net, timestamp supervision), self-training/pseudo-labeling, and data-efficient TAL? Separately: the legally-clean ways to exploit a strong teacher (Gemini) at INFERENCE for weak/pseudo-labels or as an evaluation oracle WITHOUT its outputs becoming training data for owned weights ? what does the literature say about teacher-student / distillation boundaries and label-noise handling, and what governance distinctions matter?

4) EVALUATION METHODOLOGY & RIGOR at small n: trustworthy measurement with ~6?16 videos ? bootstrap/permutation tests, paired sign-tests, leave-one-group-out vs RL00CV bias, calibration (reliability, ECE), the over-segmentation-blindness of plain tIoU, shadow/A-B evaluation, regression gates, and concrete pitfalls (label noise, annotator agreement/IAA, multiple-comparison/over-fitting traps, eval-vs-serve drift). What metrics + protocols should a small-data sports-segmentation leaderboard standardize on?

Deliverable: a structured, cited report with, per dimension, the 3?5 highest-leverage techniques that APPLY to our constrained setting (small data, local, license-clean, 1-D trajectory + cues + learned scorer), each with: what it is, why it fits our over-merge/small-data problem, the expected difficulty/payoff, and the citation. Call out which ideas are low-effort high-confidence wins vs research bets. Be explicit about anything unverified or license-incompatible."

Fetch and extract key claims from this source:

****URL:**** <https://arxiv.org/abs/2303.09055>

****Title:**** TemporalMaxer: Maximize Temporal Context with only Max Pooling for Temporal Action Localization (arXiv 2023)

****Found via:**** Single-stage TAL actionness + boundary heads, small data search

Task

1. Use WebFetch to retrieve the page content.
2. Assess source quality: primary research/institution? secondary reporting? blog/opinion? forum? unreliable?
3. Extract 2-5 FALSIFIABLE claims that bear on the research question. Each claim must:
 - be a concrete, checkable statement (not vague generalities)
 - include a direct quote from the source as support
 - be rated central/supporting/tangential to the research question
4. Note publish date if available.

If the fetch fails or the page is irrelevant/paywalled, return claims: [] and sourceQuality: "unreliable".

Structured output only.

2. assistant (2026-06-16T03:19:06.765Z)

I'll fetch the arXiv page for the TemporalMaxer paper and extract relevant claims.

3. user (2026-06-16T03:19:18.684Z)

TemporalMaxer Paper Summary

(1) Title and Authors

****Title:**** "TemporalMaxer: Maximize Temporal Context with only Max Pooling for Temporal Action Localization"

****Authors:**** Tuan N. Tang, Kwonyoung Kim, Kwanghoon Sohn

****Affiliations:**** Not explicitly listed in the provided content.

(2) Publication/Submission Date

Submitted on March 16, 2023 (arXiv:2303.09055)

(3) Core Method

TemporalMaxer employs a straightforward approach to temporal action localization. The method introduces "a basic, parameter-free, and local region operating max-pooling block" that extracts critical information from adjacent and local clip embeddings. Rather than using complex self-attention mechanisms, the approach argues that "extracted video clip features are already informative to achieve outstanding performance."

(4) Quantitative Results, Datasets, and Comparisons

The provided content does not include specific numerical results, benchmark performance metrics, dataset names, or direct comparisons to methods like ActionFormer or TriDet. These details are not present in the abstract or metadata shown.

(5) Claims About Model Complexity and Efficiency

The paper claims TemporalMaxer "outperforms other state-of-the-art methods that utilize long-term TCM such as self-attention on various TAL datasets while requiring significantly fewer parameters and computational resources."

(6) Abstract

"Temporal Action Localization (TAL) is a challenging task in video understanding that aims to identify and localize actions within a video sequence. Recent studies have emphasized the importance of applying long-term temporal context modeling (TCM) blocks to the extracted video clip features such as employing complex self-attention mechanisms. In this paper, we present the simplest method ever to address this task and argue that the extracted video clip features are already informative to achieve outstanding performance without sophisticated architectures. To this end, we introduce TemporalMaxer, which minimizes long-term temporal context modeling while maximizing information from the extracted video clip features with a basic, parameter-free, and local region operating max-pooling block. Picking out only the most critical information for adjacent and local clip embeddings, this block results in a more efficient TAL model. We demonstrate that TemporalMaxer outperforms other state-of-the-art methods that utilize long-term TCM such as self-attention on various TAL datasets while requiring significantly fewer parameters and computational resources."

4. assistant (2026-06-16T03:19:21.752Z)

The abstract page didn't include the benchmark numbers. Let me check the GitHub repo / additional context for license info and the specific datasets/numbers, which matter for the research question's license-clean constraint.

5. user (2026-06-16T03:19:29.910Z)

Temporal Action Localization Evaluation and Architecture Details

Based on the provided PDF content, I cannot extract the specific datasets, mAP numbers, parameter counts, or architectural details you requested. The document appears to be a compressed PDF with embedded figures and minimal readable text content in the stream.

The visible text fragments do not contain:

- Dataset names (THUMOS14, MultiTHUMOS, EPIC-Kitchens, ActivityNet)
- Performance metrics (mAP@tIoU values)
- Parameter counts or FLOPs comparisons
- Explicit architectural descriptions of the detection head
- Code availability or license statements

****What would help:****

To properly extract this information, I would need either:

1. A plain text or LaTeX source version of the paper
2. A decoded PDF with readable text layers
3. Direct access to the paper's main content sections (abstract, experiments, conclusion)

If you have access to the paper's source or a text-readable version, please share that instead, and I can provide the specific technical details with proper attribution and character-limited quotes.

[Binary content (application/pdf, 1.1MB) also saved to C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6a82308f-69d6-43a8-9ea4-d8eaf497253d\tool-results\webfetch-1781579965172-5ov5x2.pdf]

6. assistant (2026-06-16T03:19:32.453Z)

The PDF was saved locally. Let me read it directly to extract the numbers and architecture details.

7. user (2026-06-16T03:19:33.247Z)

pdftoppm failed: Command 'pdftoppm' not found or is in an unsafe location (current directory)

8. user (2026-06-16T03:19:38.118Z)

[notice] A new release of pip is available: 26.0.1 -> 26.1.2
 [notice] To update, run: python.exe -m pip install --upgrade pip
 ok

9. user (2026-06-16T03:19:45.179Z)

===== PAGE 1 =====

TemporalMaxer: Maximize Temporal Context with only Max Pooling
 for Temporal Action Localization

TuanN.Tang, KwonyoungKim, KwanghoonSohn*

School of Electrical and Electronic Engineering

Yonsei University

{tuanng, kyk12, khsohn}@yonsei.ac.kr

Abstract video clip features [40, 16, 60, 19, 35, 39, 46, 53]. Specifically,

the TAL model first employs pre-extracted features

Temporal Action Localization (TAL) is a challenging from a pre-trained 3D-CNN network, such as I3D [7] and

task in video understanding that aims to identify and localize TSN [52], as an input.

Then, an encoder, called backbone,

localizes actions within a video sequence. Recent studies have encoded features to latent space, and the decoder, called

emphasized the importance of applying long-term temporal head, predicts action instances as

illustrated in Fig. 1. To

context modeling (TCM) blocks to the extracted video clip better capture long-term temporal dependencies, long-term

features such as employing complex self-attention mechanisms.

In this paper, we present the simplest method to incorporate

into the backbone. Particularly, prior works [54, 61] have

addressed this task and argued that the extracted video clip features employed Graph [25], or more complex modules such as

Local-Global Temporal Encoder [40] which uses a channel

grouping strategy, or Relation-aware pyramid network [16]

introduce TemporalMaxer, which minimizes long-term temporal relations. Recently

temporal context modeling while maximizing information from

the extracted video clip features with a basic, parameter-free

attention [51] for long-term TCM [60, 35, 63, 23], resulting

free, and local region operating max-pooling block. Pick- insurprisingperformanceimprovements. ingoutonlythemostrcriticalinformationforadjacentand Long-termTCMinthebackbonecanhelpthemodelto local clip embeddings, this block results in a more effi- capture long-term temporal dependencies between frames, cientTALmodel. WedemonstratethatTemporalMaxerout- whichcanbeusefulinidentifyingcomplexactionsthatmay performs other state-of-the-art methods that utilize long- unfold over longer periods of time or heavily overlapping term TCM such as self-attention on various TAL datasets actions. Despitetheirperformanceimprovementonbench- while requiring significantly fewer parameters and com- marks, theinferencespeedandtheeffectivenessoftheex- putational resources. The code for our approach is pub- isting approaches are rarely considered. While [60] intro-

licly available at <https://github.com/TuanTNG/> ducedandpursuedtheminimalistdesign,theirbackboneis TemporalMaxer. stilllimitedtothetransformerarchitecturerequiringexpen- siveparametersandcomputations.

Recently, [58] proposed general architecture abstracted 1.Introduction from transformer. Motivated by the success of recent

approaches replacing attention module with MLP-like Temporal action localization (TAL), which aims to lo- modules[49] or Fourier Transform[26], they deemed the calize action instances in time and assign their categori- essential of those modules as token mixer which aggre- cal labels, is a challenging but essential task within the gatesinformationamongtokens. Inturn,theycametopro- field of video comprehension. Various approaches have posePoolFormer,equippedwithextremelyminimizedtoken been proposed to address this task, such as action propos- mixerwhichreplacetheexhaustingattentionmodulewith als[29],anchorwindows[36],ordenseprediction[28]. A verysimplepoolinglayer.

widely accepted notion for improving the performance of In this paper, motivated by their proposal, we focus on TALmodelsisto integrate a component capable of captur- extrememinimizationofthebackbone networkbyconcen- ing long-term temporal dependencies within the extracted trating on and maximizing the information from the ex-

*Corresponding author tractedvideo clip features in a short-term perspective rather

1
3202
raM
61
]VC.sc[
1v55090.3032:viXra
===== PAGE 2 =====
Head Head Head Head Head
Backbone Convolution Graph Transformer TemporalMaxer
MaxPooling
Pretrained 3D Pretrained 3D Pretrained 3D Pretrained 3D Pretrained 3D
CNN CNN CNN CNN CNN
Input video Input video Input video Input video Input video Clip embeddings
Common Architecture AFSD G-TAD ActionFormer Ours
Figure1.CommonarchitectureinTemporalActionLocalization(TAL)anddifferenttemporalcontextmodeling(TCM)blocks. Existing workshaveincorporatedextensiveparameters,highcomputationalcosts,andcomplexmodulesinabackbonesuch as1DConvolutional layerinAFSD[28]tocapturelocaltemporalcontext,Graph[25]inG-TAD[54]andTransformers[51]inActionFormer[60]tomodel long-termtemporalcontexts. Ourproposedmethod,termedasTemporalMaxer,forthefirsttimeexploitsthepotentialsofstrongfeatures frompretrained3DCNNbyonlyutilizingabasic,parameter-free,localoperatingMaxPoolingblock.Ourproposedmethod,thesimplest backbone ever for TAL, maintains only the most critical information on adjacent and local clip

embeddings. Combined with the large receptive field of deep networks, the whole model outperforms other works by a large margin on various datasets. than employing an exhausting long-term encoding process. tive features extracted from 3D CNN. Max Pooling [5] is The transformer architecture leading state-of-the-art performance in the machine translation task [51] and com- Our proposed method, TemporalMaxer, presents the puter vision areas [12] have inspired recent works in TAL simplest architecture ever for this task gaining much faster [60, 35, 63, 23]. From the perspective of long-term TCM, inference speed. Our finding suggests that the pre-trained the property that calculates attention weights for the long feature extractor already possess great potential and with input sequence has led to recent progress in the TAL. How- those feature, short-term TCM solely can benefit the performance, such long-term consideration comes at a price of high computation costs and the effectiveness of those approaches Extensive experiments prove the superiority and effectiveness of the proposed method, showing state-of-the-art performance in terms of both accuracy and speed for features that have not been fully exploited to date. Firstly, TAL on various challenging datasets including THUMOS the video clips exhibit a high redundancy which leads to [22], EPIC-Kitchens 100 [11], MultiTHUMOS [57], and a high similarity of the pre-extracted features as demonstrated in Fig. 2. It raises the question of the effectiveness of employing self-attention or graph methods for long-range TCM. Unlike in other domains such as machine translation task where input tokens exhibit distinctiveness, the Temporal Action Localization (TAL). Two-Stage and input clip embeddings in TAL frequently exhibit a high degree of similarity. Consequently, as depicted in Fig. 2, the videos. Two-Stage methods first generate possible action self-attention recently employed in TAL tend to average the proposals and classify them into actions. The proposals are embeddings of clips within the attention scope, losing temporal local minute changes by redundant similar frame representations [4, 54], under the long-term temporal context modeling. We argue that Transformers [53, 8, 46]. Single-stage TAL performs that only certain information within clip embeddings is relevant to the action context of TAL, whereas the remainder pass, without using a separate proposal generation step. The of the information is similar across adjacent clips. Their pioneering work [40] developed anchor-based single-stage fore, an optimal TCM must be capable of preserving the most discriminative features of clip embeddings that carry stage object detector [42, 33]. Meanwhile, [28] proposed the essential information. an anchor-free single-stage model with a saliency-based regression module. To this end, we aim to propose a simple yet effective TCM refinement module. in a straightforward manner. We argue that the TCM can Long-term Temporal Context Modeling (TCM). Retain the simplest architecture while maximizing information. cent studies have emphasized the necessity of long-term

===== PAGE 3 =====

Former [60] as the baseline for comparison. Similar to ActionFormer, TemporalMaxer follows a minimalistic design of sequence labeling where every moment is classified, and their corresponding action boundaries are regressed. The

Attention Score main difference is that we avoid exhausting attention be-

5.17

tween clips from long-term time frames which can unintentionally

1.06

5 5.5 5 AvgPooling tionally flatten minute information among the crowd of sim-

3.03

1.2 1 1 ilar frames, but keep the minute information in short-term

3 3.1 3 MaxPooling 5.5 manner.

1.2

Clip 1 2 3

3.1

3. Method

(a) Simplified structure of Attention and Max Pooling module

3.1. Problem Statement

0 1.0 0 0.15

Temporal Action Localization. Assume that an

10 0.9 10 0.12 untrimmed video X can be represented by a set of feature vectors $X = \{x_1, x_2, \dots, x_T\}$, where the number of dis-

20 0.8 20 1 2 T

0.09 create time steps $t = \{1, 2, \dots, T\}$ may vary depending on

30 0.7 30 the length of the video. The feature vector x_t is extracted

0.06 from a pre-trained 3D convolutional network and represents

40 0.6 40 a video clip at a specific moment t . The aim of TAL is

to predict a set of action instances $y = \{y_1, y_2, \dots, y_T\}$

50 0.5 50 0.03 1 2 N

based on the input video sequence X , where N is the number of

(b) Cosine similarity of (c) Attention Score

for action instances in X . Each action instance y_i con-

n

sists of (s_i, e_i, a_i) where s_i, e_i and a_i are starting time,

Figure 2. Comparison of effectiveness of Max Pooling over $n \times n \times n \times n \times n$

ending time, and associated action label a respectively,

Transformer [51]. (a) When the input clip embeddings are highly n

similar, given that the value features V are highly similar and the $s_n \in [1, T]$, $e_n \in [1, T]$, $s_n < e_n$, and the action label

attention score is distributed equally, the Transformer tends to average out the clip embeddings as done in Average Pooling. In con-

trast, Max Pooling can retain the most crucial information about

3.2. Temporal Maxer

adjacent clip embeddings and further remove redundancy in the Action Representation. We follow

the anchor-free video. (b) Cosine similarity matrix (bottom-left) of value features single-stage representation [28, 60] for an action instance.

in self-attention in ActionFormer [60] where the red boxes exhibit

Each moment is classified into the background or one of C

the action intervals. (c) Attention score (bottom-right) after train-

categories, and regressed the onset and offset based on the

in the ActionFormer [60].

current time step of that moment. Consequently, the predic-

tion in TAL is formulated as a sequence labeling problem.

TCM to improve model performance. Long-term TCM $X = \{x_1, x_2, \dots, x_T\}$??? = (cid:110) ??, ??, ..., ??

(cid:111) (1)

1 2 T 1 2 T

helps the model capture long-term temporal dependencies

between frames, which can be useful in identifying intricate At the time step t , the output $?? = (os, oe, c$

) is defined as

action that spans over extended time frames or heavily over- $t \times t \times t$

following:

lapping actions. Prior work has addressed long-term TCM

using Relation-aware pyramid Network [16], Multi-Stage $os > 0$ and $oe > 0$ represent the

temporal intervals

$t \times t$

CNN [14], Temporal Context Aggregation Network [40].

between the current time step and the onset and offset

Recent works [54, 61] employ Graph [25] where each video of a given moment, respectively.

clip feature represents a node in a graph as long-term TCM.

? Given action categories, the action probability can

More recently, Transformer [51] demonstrates an outstanding

be considered as a set of c_i which is a probability for

ing capacity to capture long-ranged dependency of the input t

action i at t , where $1 \leq i \leq C$.

sequence in the machine translation tasks. Thus, it is a nat-

ural fit for Temporal Action Localization (TAL) where each

The predicted action instance at time step t can be retrieved

video clip embedding represents a token. Therefore, recent from $?? = (os, oe, c)$ by:

$t \ t \ t \ t$

studies [60, 35, 63, 23] have employed Transformers as a

$a = \arg\max(c)$, $s = t \cdot os$, $e = t + oe$ (2)

long-term TCM. $t \ t \ t \ t \ t$

Our approach, TemporalMaxer, belongs to the single- Architecture overview. The overall

architecture is de-

stage TAL model that utilizes a state-of-the-art Action- picted in Fig. 3. Our proposed method

aims to learn to

===== PAGE 4 =====

TemporalMaxer Cls + Reg Head

.....

TemporalMaxer

Max Pooling

..... Cls + Reg Head

Sliding Direction TemporalMaxer

... Cls + Reg Head

TemporalMaxer

Feature projection layer

Pre-trained 3D CNN

... ..

Figure 3. Overview of TemporalMaxer. The proposed method utilizes Max Pooling as a Temporal Context Modeling block

applied between

temporal feature pyramid levels to maximize informative features of high similarity clip embedding. Specifically,

it first extracts features of

every clip using pre-trained 3D CNN. After that, the backbone encodes clip features to form a multi-

scale feature pyramid. The backbone

consists of 1D convolutional layers and TemporalMaxer layers. Finally, a lightweight classification and regression

head decodes the feature

pyramid to action candidates for every input moment.

label every input moment by $f(X)$?? with f is a deep encoded by TemporalMaxer:

learning model. f follows an encoder-decoder design and

$Z_l = \text{TemporalMaxer}(Z_l; 1)$. (4)

can be decomposed as d . The encoder here is the back-

bone, and the decoder is the classification and regression

Here: TemporalMaxer is Max Pooling and employed with

X head n . to e la : t X nt ? vec Z tor le Z ar , n a t n o d e d nc : od Z e t ? he ? i ?

np l u e t ar v n i s d e t o p f e r a e t d u i r c e t stride 2, Z_l ? R $2l$ T ? $1 \times D$, $2 \leq l \leq$

L . It is worth not-

ing that ActionFormer [60] employs Transformer [51] as a

labels for every input moment. To effectively capture ac-

TCM block where each clip feature at the moment t repre-

sents a token, our proposed method adopts only Max Pool-

a multi-scale feature pyramid representation, which is de-

ing [5] as TCM block.

noted as $Z = \{Z_1, Z_2, \dots, Z_L\}$.

Decoder Design The decoder d learns to predict

Encoder design The input feature X is first en- sequence labeling, $?? = (cid:110) ?? , ?? , \dots, ??$

(cid:111) , for ev-

coded into multi-scale temporal feature pyramid $Z = 1 \ 2 \ T$

$\{Z_1, Z_2, \dots, Z_L\}$ using encoder e . The encoder e simply every moment using multi-scale feature

pyramid $Z =$

$\{Z_1, Z_2, \dots, Z_L\}$. The decoder adopts a lightweight con-

tainst two 1D convolutional neural network layers as fea-

volutional neural network and consists of classification and regression heads. Formally, the two heads are defined as:

textModeling(TCM) block to produce feature pyramid Z . Formally, the feature projection layers are described as:

$$C = F(E(E(Z_1))) \quad (5)$$

$$l = c \times 4 \times 3$$

$$O = \text{ReLU}(F(E(E(Z_1)))) \quad (6)$$

$$X = E(E(\text{Concat}(X))) \quad (3) \quad l = o \times 6 \times 5$$

$$p = 2 \times 1$$

Here, $Z_l \in \mathbb{R}^{2l \times T \times 1 \times c}$ is the latent feature of level l . The input video feature sequence $X = \{x_1, x_2, \dots, x_T\}$, $l, C, l = \{c_0, c_{2l \times 1}, \dots, c_T\} \in \mathbb{R}^{2l \times T \times 1 \times c}$ denotes the where $x_i \in \mathbb{R}^{1 \times D_{in}}$, is first concatenated in the first dimension and then fed into two feature projection modules $\{(os_0, oe_0, c_0), (os_{2l \times 1}, oe_{2l \times 1}, c_{2l \times 1}), \dots, (os_T, oe_T, c_T)\} \in \mathbb{R}^{2l \times T \times 1 \times 2}$ is the on- E , and E in equation 3, resulting in the projected feature set and offset prediction of input moment $\{0, 2l \times 1, \dots, T\}$.

$$l = 2$$

ture $X \in \mathbb{R}^{T \times D}$ with D -dimensional feature space. Each E denotes the 1D convolution followed by Layer Normalization.

projection module comprises one 1D-convolutional neuralization and ReLU activation function. F and F are both

$$c = 0$$

network layer, followed by Layer Normalization [3], and 1D convolution. Note that all the weights of the decoder

ReLU [1]. We simply assign $Z_1 = X$ as the first feature are shared between the different features in the multi-scale

$$p = 1$$

in Z . Finally, the multi-scale temporal feature pyramid Z is feature pyramid Z .

===== PAGE 5 =====

Learning Objective. The model predicts $?? = 4.1$. Results on THUMOS14

t

(os, oe, c) for every moment of the input X . Following the

t t t Dataset. THUMOS14 dataset [22] contains 200 validation videos and 213 testing videos with 20 action classes. employed to supervise classification and regression outputs

Following previous work [29, 30, 54, 62, 60], we trained respectively. The overall loss function is defined as:

the model using validation videos and measured the performance on testing videos.

$$L = (c_{id:88}) (L + 1 \times L) / T \quad (7)$$

total cls ct reg +

Feature Extraction. Following [60, 62], we extract the

t

features of THUMOS14 dataset using two-stream I3D [7] where L denotes regression loss and is applied only reg pre-trained on Kinetics [24]. 16 consecutive frames are fed when the indicator function, 1 , indicates that the current ct into I3D pre-trained network with a sliding window of stride timestep is a positive sample. T is the number of positive + 4. The extracted feature is collected after the last fully con-

samples. L is CWay classification loss. The loss function is connected layer and has 1024-D feature space. After that, the L is applied to all levels on the output of multi-scale total two-stream features are further concatenated (2048-D) and feature pyramid Z and averaged across all video samples utilized as the input of the model.

during training.

Results. We compare the performance evaluated on the THUMOS14 dataset [22] with state-of-the-art methods.

4. Experimental Results

TemporalMaxer demonstrates remarkable performance.

In this section, we show that our proposed method, TemporalMaxer, achieving an average mAP of 67.7% mAP, out-

performs, demonstrates the outstanding results achieved performing all previous approaches, both single-stage, and

across a variety of challenging datasets, namely THUMOS two-stage methods, by a significant margin, with a 1.1% increase in mAP.

[22], EPIC-Kitchens 100 [11], MultiTHUMOS [57], and

creases in mAP at IoU=0.4. Especially, TemporalMaxer sur-

MUSES [34]. These datasets are recognized as standard passes all recent methods that utilize long-term TCMB blocks

10. assistant (2026-06-16T03:19:48.935Z)

I have the key data. Let me get the rest of the results tables to capture the efficiency/parameter numbers and comparisons to ActionFormer.

11. user (2026-06-16T03:19:54.379Z)

===== PAGE 5 =====

Learning Objective. The model predicts ?? = 4.1. Results on THUMOS14

t

(os, oe, c) for every moment of the input X . Following the THUMOS14 dataset [22] contains 200 validation videos and 213 testing videos with 20 action classes.

employed to supervise classification and regression outputs

Following previous work [29, 30, 54, 62, 60], we trained

respectively. The overall loss function is defined as:

the model using validation videos and measured the performance

$L = (cid:88) (L + 1 L) / T (7)$ mance on testing videos.

total cls ct reg +

Feature Extraction. Following [60, 62], we extract the

t

features of THUMOS14 dataset using two-stream I3D [7]

where L denotes regression loss and is applied only

reg pre-trained on Kinetics [24]. 16 consecutive frames are fed

when the indicator function, 1 , indicates that the current

ct into I3D pre-trained network with a sliding window of stride

time step is a positive sample. T is the number of positive

+ 4. The extracted feature is collected after the last fully con-

samples. L is CWay classification loss. The loss function

is connected layer and has 1024-D feature space. After that, the

L is applied to all levels on the output of multi-scale

total two-stream features are further concatenated (2048-D) and

feature pyramid Z and averaged across all video samples

utilized as the input of the model.

during training.

Results. We compare the performance evaluated on the THUMOS14 dataset [22] with state-of-the-art methods.

4. Experimental Results

TemporalMaxer demonstrates remarkable performance.

In this section, we show that our proposed method, TemporalMaxer, achieving an average mAP of 67.7% mAP, out-

performs, demonstrates the outstanding results achieved performing all previous approaches,

both single-stage, and across a variety of challenging datasets, namely THUMOS two-stage methods, by a significant margin, with a 1.1% increase in mAP at $tIoU=0.4$. Especially, Temporal Maxersur-MUSES [34]. These datasets are recognized as standard passes all recent methods that utilize long-term TCM blocks benchmarks in the Temporal Action Localization task. Our including self-attention such as TadTR [35], HTNet [23], approach surpasses the state-of-the-art baseline, Action-TAGS [38], GLFormer [19], ActionFormer [60], or GraphFormer [60], in each dataset, showcasing its superior performance compared to other works. including Local-Global Temporal Encoder [40].

Evaluation Metric. We employ a widely-used evaluation metric for TAL known as the mean average precision (mAP) calculated at various temporal intersections exhibits outstanding performance but also is efficient in over union ($tIoU$). $tIoU$ is the intersection over union between two temporal windows, i.e., the 1D Jaccard index. takes 50ms on average to fully process an entire video on We report the mAP scores for all action categories based THUMOS. It is 1.6x faster than the ActionFormer baseline on the given $tIoU$ thresholds, and further report an averaged and 3.9x faster than TadTR. However, in section 4.5, we mAP value across all $tIoU$ thresholds. show that our model only takes 10.4ms for forward time. It

Training Details. To ensure a fair and unbiased comparison, we employed the experimental setup of the robust baseline model, ActionFormer [60]. This setup included various components such as decoder design d, and 8.0x faster than ActionFormer, respectively. non-maximum suppression (NMS) hyper-parameters in the

4.2. Results on EPIC-Kitchens100

post-processing stage, data augmentation, learning rate, feature extraction, and the number of feature pyramid levels. Dataset. The EPIC-Kitchens100 dataset [11] is a comprehensive collection of ego-centric action videos, featuring Transformer block in ActionFormer with the Max Pooling 100 hours of footage from 700 sessions that document cooking activities in a variety of kitchens. Additionally, EPIC-3 for all TCM blocks. The subsequent ablation will thoroughly analyze the effects of varying kernel sizes. During hours and more than ten times larger in terms of action instances (averaging 128 per video) when compared to THUMOS14. These videos are recorded from a first-person perspective, resulting in insignificant camera motion, and represent on the EPIC-Kitchens100 dataset, and approximately sent a novel challenge for TAL research. 45 minutes on MUSES. Additionally, ModelEMA [21] and Feature Extraction. Following previous work [60, 29, gradient clipping techniques are employed, consistent with [54], we extract the video's feature using SlowFast network those used in [60], to promote training stability. [15] pre-trained on EPIC Kitchens [11]. We utilized a

tIoU?

Type Model Feature time(ms)?

0.3 0.4 0.5 0.6 0.7 Avg.

BMN[29] TSN[52] 56.0 47.4 38.8 29.7 20.5 38.5 483*

DBG[27] TSN[52] 57.8 49.4 39.8 30.2 21.7 39.8 ?

G-TAD[54] TSN[52] 54.5 47.6 40.3 30.8 23.4 39.3 4440*

BC-GNN[4] TSN[52] 57.1 49.1 40.4 31.2 23.1 40.2 ?

TAL-MR[62] I3D[7] 53.9 50.7 45.4 38.0 28.5 43.3 >644*

P-GCN[59] I3D[7] 63.6 57.8 49.1 ? ? ? 7298*

P-GCN[59]+TSP[2] R(2+1)1D[50] 69.1 63.3 53.5 40.4 26.0 50.5 ?

Two-Stage

TSA-Net[17] P3D[41] 61.2 55.9 46.9 36.1 25.2 45.1 ?

MUSES[34] I3D[7] 68.9 64.0 56.9 46.3 31.0 53.4 2101*

TCANet[40] TSN[52] 60.6 53.2 44.6 36.8 26.7 44.3 ?

BMN-CSA[45] TSN[52] 64.4 58.0 49.2 38.2 27.8 47.7 ?

ContextLoc[65] I3D[7] 68.3 63.8 54.3 41.8 26.2 50.9 ?

VSGN[61] TSN[52] 66.7 60.4 52.4 41.0 30.4 50.2 ?

RTD-Net[46] I3D[7] 68.3 62.3 51.9 38.8 23.7 49.0 >211*

Disentangle[66] I3D[7] 72.1 65.9 57.0 44.2 28.5 53.5 ?

SAC[55] I3D[7] 69.3 64.8 57.6 47.0 31.5 54.0 ?

A²Net[56] I3D[7] 58.6 54.1 45.5 32.5 17.2 41.6 1554*

GTAN[36] P3D[41] 57.8 47.2 38.8 ? ? ? ?

PBRNet[32] I3D[7] 58.5 54.6 51.3 41.8 29.5 ? ?

Single-Stage AFSD[28] I3D[7] 67.3 62.4 55.5 43.7 31.1 52.0 3245*

TAGS[38] I3D[7] 68.6 63.8 57.0 46.3 31.8 52.8 ?

HTNet[23] I3D[7] 71.2 67.2 61.5 51.0 39.3 58.0 ?

TadTR[35] I3D[7] 74.8 69.1 60.1 46.6 32.8 56.7 195*

GLFormer[19] I3D[7] 75.9 72.6 67.2 57.2 41.8 62.9 ?

AMNet[35] I3D[7] 76.7 73.1 66.8 57.2 42.7 63.3 ?

ActionFormer[60] I3D[7] 82.1 77.8 71.0 59.4 43.9 66.8 80

ActionFormer[60]+GAP[37] I3D[7] 82.3 ? 71.4 ? 44.2 66.9 >80

Our(TemporalMaxer) I3D[7] 82.8 78.9 71.8 60.5 44.7 67.7 50

Table 1. The results obtained on the THUMOS14 dataset [22] are presented for various IoU thresholds, with the average mAP calculated

in the range [0.3:0.7:0.1]. The top-performing results are highlighted in bold.

The time (ms) is the average inference time for one video,

without extracting features from 3D CNN and including the post-processing step, such as NMS. We measure the inference time using

asingle GeForce GTX 1080 Ti GPU. Results indicated with * are taken from [35] which are measured using Tesla P100 GPU, a much

more powerful GPU than the 1080 Ti. In comparison to early works, including both one-stage and two-stage methods, and those utilizing

long-term TCM, TemporalMaxer achieves superior performance in both mAP and inference speed.

Frame input sequence with a stride of 16 to generate a set of video hours with 25 action classes, designed to facilitate

2304-D features. These features were then fed as input to multi-shot analyses making the dataset challenging.

our model.

Result. Tab. 2 shows our results. TemporalMaxer

Feature Extraction. We directly employ the pre-

demonstrates notable performance on the EPIC-Kitchens

extracted feature provided by [34]. The feature is extracted

100 dataset, achieving an average mAP of 24.5% and 22.8%

using a pre-trained I3D network [7] on the Kinetics dataset

for verb and noun, respectively. The superiority of our ap-

[24] using only RGB stream, resulting in 1024-D feature

proach is further confirmed by a large margin over a strong

space for a video clip embedding.

and robust baseline, ActionFormer [60], with an average

improvement of 1.0% mAP for verb and 0.9% mAP for

noun. Again, TemporalMaxer outperforms other methods

Result Tab. 3 shows the results on MUSES dataset. Our

that utilizes long-term TCM including self-attention [60] or

method significantly outperforms other works that employ

Graph [54]. These results provide empirical evidence of

long-term TCM such as G-TAD [54], Ag-Trans [63], and the effectiveness of the simple backbone in advancing the ActionFormer [60] at every tIoU threshold. Notably, TemporalMaxer improves 1.0 mAP at tIoU=0.7. On average, we achieve 27.2 mAP, which is 1.0 mAP higher than the previous approaches, demonstrating the robustness of our Dataset. The MUSES dataset [34] is a collection of method. It is worth noting that we implemented the ActionFormer [60] on the MUSES dataset using the code provided by the authors.

4.3. Results on MUSES

previous approaches, demonstrating the robustness of our Dataset. The MUSES dataset [34] is a collection of method. It is worth noting that we implemented the ActionFormer [60] on the MUSES dataset using the code provided by the authors.

===== PAGE 7 =====

tIoU

Task Method

0.1 0.2 0.3 0.4 0.5 Avg

BMN[29,11] 10.8 9.8 8.4 7.1 5.6 8.4

G-TAD[54] 12.1 11.0 9.4 8.1 6.5 9.4

Verb

ActionFormer[60] 26.6 25.4 24.2 22.3 19.1 23.5

Our(TemporalMaxer) 27.8 26.6 25.3 23.1 19.9 24.5

BMN[29,11] 10.3 8.3 6.2 4.5 3.4 6.5

G-TAD[54] 11.0 10.0 8.6 7.0 5.4 8.4

Noun

ActionFormer[60] 25.2 24.1 22.7 20.5 17.0 21.9

Our(TemporalMaxer) 26.3 25.2 23.5 21.3 17.6 22.8

Table 2. The performance of four proposed methods on the EPIC-Kitchens100 dataset [11] is evaluated using various tIoU thresholds. The average mAP is reported over a range of tIoU thresholds [0.1:0.5:0.1]. The top-performing methods are highlighted in bold. Our proposed method outperforms the other methods significantly.

tIoU tIoU

Method Method

0.3 0.5 0.7 Avg 0.2 0.5 0.7 Avg

BU-TAL[62] 12.9 9.2 5.9 9.4 PDAN[10] ? ? ? 17.3

G-TAD[54] 19.1 11.1 4.7 11.4 MLAD[48] ? ? ? 14.2

P-GCN[59] 19.9 13.1 5.4 13.0 MS-TCT[9] ? ? ? 16.2

MUSES[34] 25.9 18.9 10.6 18.6 PointTAD[47] 39.7 24.9 12.0 23.5

Ag-Trans[63] 24.8 19.4 10.9 18.6 ActionFormer[60] 46.4 32.4 15.0 28.6

ActionFormer[60] 35.9 26.9 15.2 26.2 Our(TemporalMaxer) 47.5 33.4 17.4 29.9

Our(TemporalMaxer) 36.7 27.8 16.2 27.2 Table 4. Comparison with the state-of-the-art methods on the MultiTHUMOS dataset. We report the results at different tIoU thresholds [0.3, 0.5, 0.7] and average mAP in [0.1:0.9:0.1].

and the average mAP in [0.3:0.1:0.7] on MUSES dataset [34]. All methods used the same I3D features. Our method outperforms concurrent works by a large margin.

mance on the MultiTHUMOS dataset [57] with recent state-of-the-art methods. Specifically, our method surpasses the prior work [47] that utilizes Transformer as feature en-

4.4. Results on MultiTHUMOS

coding by a large margin, 6.4% mAP on average. More-

Dataset. The MultiTHUMOS dataset [57] is a densely

over, TemporalMaxer improves the robust baseline, Action-labeled extension of THUMOS14, consisting of 413 sports videos with 65 distinct action classes. The dataset presents average. It should be noted that we utilized the code pro-

a significant increase in the average number of distinctive video by the authors to implement ActionFormer [60] on action categories per video, compared to THUMOS14. As the MultiTHUMOS dataset, and the results [10, 48, 9] are

such, it poses a greater challenge for TAL than THUMOS. taken from [47].

While MultiTHUMOS are being used in action detection

4.5. Ablation Study

benchmark [10, 9], a novel approach for action detection, PointTAD[47], utilizes the TAL evaluation metric to assess the completeness of predicted action instances. Given that the effectiveness of TemporalMaxer. To better understand what is TAL and action detection share the same setting in terms of the effective component for TCM, we gradually replace the input features and annotations, we evaluate the performance of Max Pooling with other blocks such as convolution, sub-sampling, and Average Pooling. Furthermore, we evaluate state-of-the-art action detection methods [10, 48, 9, 47], and numerous kernel sizes of Max Pooling. Note that all experimental strong baseline ActionFormer [60].

12. assistant (2026-06-16T03:19:58.655Z)

I now have comprehensive data including the key efficiency claims. Let me extract the ablation page to capture the comparison of max pooling vs other TCM blocks, which is highly relevant to the over-segmentation question.

13. user (2026-06-16T03:20:03.219Z)

tIoU

TCM GMACs? #params(M)? time(ms)? backbone time(ms)?

0.5 0.7 Avg.

Conv[28](Our Impl) 62.8 37.1 59.4 45.6 30.5 16.3 9.0

Subsampling 64.3 37.7 61.0 16.2 7.1 10.4 2.5

Average Pooling 66.1 39.4 63.2 16.4 7.1 10.4 2.5

Transformer[51] 71.0 43.9 66.8 45.3 29.3 30.5 20.1

TemporalMaxer 71.8 44.7 67.7 16.4 7.1 10.4 2.5

Table 5. Ablation studies about different TCM blocks on THUMOS14.

Inference times are measured using an input video with 2304 clip embeddings, a 5-minute video, on a GeForce GTX 1080 Ti GPU without post-processing (NMS) and pre-extracted features step.

from 3D-CNN for the TAL task. PoolFormer retains the global information of adjacent clips is reduced. That is why

structure of Transformer such as FFN [43, 44], residual Average Pooling decreases by 3.6% mAP compared with

connection [18], and Layer Normalization [3] because the Transformer.

tokens have to be encoded by the PoolFormer itself. However, the ablation study with Average Pooling suggests that

ever, in TAL the features from 3D CNN have already been the most important information of clip embeddings should pre-extracted and contain useful information. Therefore, be maintained.

For that reason, we employ Max Pooling as

we posit that there will be a straightforward block that are

TCM. The result is provided in the last row of Tab. 5. Max

more efficient than Transformer/PoolFormer, does not re-

quire much computational as well as parameters, and be and significantly outperforms the strong

and robust base-

able to effectively exploit the pre-extracted features. line, ActionFormer. To clarify,

TemporalMaxer effectively

highlights only critical information from nearby clips and

Our ablation study starts to employ a 1-D convolution

discards the less important information. This result suggests

module as an alternative to Transformer[51] for TCM block

that the feature from pretrained 3D CNN are informative

to see how many mAP drop. The result is reported in the

and can be effectively utilized for the TAL model without

first row of Tab. 5. As expected, the convolution layer

complex modules like prior works.

decreases by 7.4% mAP compared with the ActionFormer

Our method, TemporalMaxer, results in the simplest

baseline. This deduction can be explained for two reasons.

model ever for TAL task that contains minimalist parameters. First, given the video clip features that are informative but expensive and computational cost for the TAL model. TemporalMaxer is highly similar, and the convolution weight is fixed after training, consequently the convolution operation cannot retain the most informative features of local clip embeddings. Second, the convolution layer introduces the most parameters when comparing only the backbone time, our proposed method only takes 2.5 ms which is incredibly 8.0x faster than ActionFormer backbone [63], 20.1ms. trained 3D CNN, the model should not contain too many parameters, which may lead to overfitting and thus less generalization. To further concretize our thought, we replace the average mAPs are 67.7, 67.1, 66.8, 65.7, respectively. Our Transformer with a parameter-free operation, subsampling model achieved the highest performance with a kernel size of 3, while the lowest performance was observed with a kernel size of 6. This decrease in performance can be attributed to the corresponding loss of information during training. The results obtained for kernel sizes 3, 4, and 5 are not very sensitive to the kernel size. This suggests that these kernel contain many parameters, and the features from pre-trained sizes can effectively capture the relevant temporal information for the task at hand. However, subsampling is prone to losing the most crucial information as only half of the video clip embeddings are kept after a TCMB block. Thus, we replace Transformer with Average Pooling with kernel size 3 and stride 2. As expected, Average Pooling improves the result by 2.2 % approach, TemporalMaxer for temporal action localization mAP on average compared with the subsampling. It is task. To minimize the structure, we explore the way to because this operation does not drop any clip embedding simply maximize the underlying information in the video which helps to retain the crucial information. However, the clip features from pre-trained 3D-CNN. To this end, with a Average Pooling averages the nearby features, thus the crucial, non-parametric and temporally local operating max-

14. assistant (2026-06-16T03:20:08.231Z)

I now have complete data: datasets, exact mAP numbers, efficiency comparisons, parameter counts, ablation results, license/code availability, and architecture details. I have everything needed to assess the source and extract falsifiable claims.

Source assessment: This is a primary peer-reviewed research paper (CVPR 2023 workshop / arXiv) from Yonsei University. Code is publicly available on GitHub. The method is highly relevant to Dimension 1 (windowing/boundaries, single-stage TAL, actionness+boundary heads) and Dimension 3 (small-data ? the parameter-free design explicitly argues fewer params reduce overfitting).

15. user (2026-06-16T03:20:23.957Z)

Structured output provided successfully

16. assistant (2026-06-16T03:20:34.799Z)

Extracted and verified the claims from the TemporalMaxer paper (arXiv:2303.09055, Tang/Kim/Sohn, Yonsei University, submitted 2023-03-16). Source is a primary research paper with publicly released code.

The most load-bearing findings for the research question (closing the over-merge gap on tiny data):

- TemporalMaxer is ActionFormer with the self-attention block swapped for a parameter-free MaxPooling block ? identical anchor-free actionness + onset/offset-regression head, Focal + DIOU loss. This is directly the "actionness + boundary heads" single-stage TAL the project already cites, in its leanest form.
- It beats ActionFormer on THUMOS14 (67.7 vs 66.8 avg mAP) with ~4x fewer params (7.1M vs 29.3M) and 2.8x fewer GMACs ? and wins on all four datasets tested.
- The ablation is the key small-data signal: learned 1-D conv TCM is the worst (59.4, overfits via too many params); parameter-free subsampling/avg-pool/max-pool all beat it. The authors explicitly argue more learned temporal params hurt generalization on informative-but-redundant features.
- They diagnose self-attention as an over-merging operation that averages similar adjacent clips and loses local minute changes ? the same merge-vs-preserve tension as the project's mega-window failure.

Note for the synthesis: the paper's claims rest on standard benchmarks (THUMOS14, EPIC-Kitchens-100, MultiTHUMOS, MUSES), not on a 1-D shuttle-trajectory or tiny-n (6?16 video) setting. The "fewer params generalize better" argument is suggestive for small data but is not measured at the project's data scale ? that transfer is unverified. The MIT/Apache license status was not stated on the paper page; only "code publicly available" on GitHub is confirmed, so the license itself needs separate verification against the project's MIT/Apache/BSD-only constraint.