

claude-windows / 6cee04e1-f4f1-4593-86c9-2e3ca3473efd

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\6cee04e1-f4f1-4593-86c9-2e3ca3473efd\subagents\workflows\wf_e3b53aed-f4e\agent-a53f4ee74c06ce749.jsonl`
- SHA-256: `fdef145308c5aa08c30d48f2019510b9becc7e2bfdc4a48ab57b78e3bbc1edd8`
- Source modified: `2026-06-18T02:03:38+00:00`
- Imported at: `2026-07-05T16:48:26+00:00`
- project: `wf_e3b53aed-f4e`
- session_id: `6cee04e1-f4f1-4593-86c9-2e3ca3473efd`
```

Transcript

1. user (2026-06-18T02:00:03.302Z)

Review a bug-fix PR in badminton-highlight-indexer. Work READ-ONLY in this git worktree (do NOT modify): C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/scripts-unwrap

First run ``git -C "C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/scripts-unwrap" diff origin/master..HEAD`` to see the exact change, then read the touched files + related code/tests. PR #213 fixes two standalone scripts (scripts/run_phase3_eval.py, scripts/run_process_and_stitch.py) that called `len()/truthiness` on a raw segmenter Result (Success|Failure dataclass, no `__len__`) ? crashing under the default 'motion' segmenter. Adds backend/results.unwrap_or_failure(result) -> (value, failure) (DRY of the backend/main.py inline unwrap) + tests/test_results.py.

Be adversarial ? try to find a real reason NOT to merge. Concrete, verifiable findings only (cite file:line); no style noise. severity must-fix only for a genuine bug/regression/break. blocks_merge=true ONLY on a real must-fix. Return the structured object.

LENS ? scope & convention: exactly the right files touched? No unintended behavior/product/licensing change? Matches surrounding style? Commit/PR claims accurate? For #213, is adding a helper to backend/results.py (vs inlining) reasonable and non-disruptive to main.py's existing inline sites?

2. assistant (2026-06-18T02:00:05.433Z)

I'll review this bug-fix PR. Let me start by examining the exact change.

3. user (2026-06-18T02:00:07.044Z)

```
diff --git a/backend/results.py b/backend/results.py
index d6b1ea4..75120c5 100644
--- a/backend/results.py
+++ b/backend/results.py
@@ -18,7 +18,7 @@ This module has no dependencies on the rest of the app.
     from __future__ import annotations

     from dataclasses import dataclass, field
     -from typing import Generic, TypeVar, List, Optional, Any, Union
     +from typing import Generic, TypeVar, List, Optional, Any, Tuple, Union

     T = TypeVar("T")

@@ -58,3 +58,21 @@ class Failure:
     # Result[ValueType] means Success[ValueType] | Failure
     # (the Failure "error" type is fixed for this contract)
```

```

Result = Union[Success[T], Failure]
+
+
+def unwrap_or_failure(result: Any) -> Tuple[Optional[Any], Optional[Failure]]:
+    """Split a segmenter/provider return into ``(value, failure)`` so a caller branches once.
+
+    - ``Failure`` -> ``(None, failure)``
+    - ``Success(value=v)`` -> ``(v, None)``
+    - a bare value (legacy code that predates this contract and still returns a raw list)
+    -> ``(value, None)``
+
+    This is the same isinstance/``.value`` unwrap that ``backend/main.py`` does inline; callers
+    that forget it (e.g. ``len(process_video(...))``) crash with ``TypeError: ... has no
+    len()``
+    on a ``Success`` dataclass. ``result`` is typed ``Any`` because the contract is
+    heterogeneous
+    (Success | Failure | legacy bare value)."""
+    if isinstance(result, Failure):
+        return None, result
+    value = result.value if hasattr(result, "value") else result
+    return value, None
diff --git a/scripts/run_phase3_eval.py b/scripts/run_phase3_eval.py
index 56e21ca..dd1263e 100644
--- a/scripts/run_phase3_eval.py
+++ b/scripts/run_phase3_eval.py
@@ -34,6 +34,7 @@ from backend.config import load_config
from backend.utils.hashing import compute_video_id as get_file_md5 # canonical video_id
# Importing the package registers all segmenters (motion/gemini/yolo_hybrid/...).
from backend.pipeline.segmenters import segmenter_registry
+from backend.results import unwrap_or_failure
from backend.eval.annotations import load_annotations
from backend.eval.harness import evaluate, format_report_text, report_to_dict
from backend.eval import batch
@@ -142,7 +143,15 @@ def main(argv=None) -> int:
    print(f"[proc] {vid}: segmenting with {args.segmenter}...")
    db.add_video(video_id, video_path, "processed", [])
    t0 = time.time()
-    segs = segmenter.process_video(video_id, video_path,
max_frames=args.max_frames)
+    # process_video returns a Result (Success|Failure) for some segmenters
+    # (e.g. the default `motion`); unwrap before len() so this doesn't crash with
+    # TypeError on a Success dataclass ? mirrors the unwrap in backend/main.py.
+    segs, failure = unwrap_or_failure(
+        segmenter.process_video(video_id, video_path, max_frames=args.max_frames))
+    if failure is not None:
+        print(f"        -> FAILED: {failure.message}")
+        skipped.append((vid, f"segment failed: {failure.message}"))
+        continue
    print(f"        -> {len(segs)} segments in {time.time() - t0:.0f}s")

    gt = load_annotations(csv_path)
diff --git a/scripts/run_process_and_stitch.py b/scripts/run_process_and_stitch.py
index b8ffef6..ea7c2de 100644
--- a/scripts/run_process_and_stitch.py
+++ b/scripts/run_process_and_stitch.py
@@ -16,6 +16,7 @@ from backend.config import load_config
from backend.utils.hashing import compute_video_id
from backend.pipeline.segmenters import segmenter_registry
from backend.pipeline.stitcher.compiler import VideoCompiler
+from backend.results import unwrap_or_failure
from backend.tools.export import format_rally_summary

logger = logging.getLogger(__name__)
@@ -78,7 +79,13 @@ def main():
    return

```

```

        print(f"Using segmenter: {seg_name}", flush=True)
        segmenter = segmenter_cls(db, cfg)
-       segments = segmenter.process_video(video_id, video_path)
+       # process_video returns a Result (Success|Failure) for some segmenters (e.g. `motion`);
+       # unwrap before len()/truthiness so this doesn't crash with TypeError on a Success
+       # dataclass (and so the empty-guard below actually sees the list) ? mirrors
backend/main.py.
+       segments, failure = unwrap_or_failure(segmenter.process_video(video_id, video_path))
+       if failure is not None:
+           print(f"[FAIL] Segmentation failed: {failure.message}")
+           return
        print(f"Successfully indexed {len(segments)} play segments.", flush=True)

    if not segments:
diff --git a/tests/test_results.py b/tests/test_results.py
new file mode 100644
index 0000000..5b50767
--- /dev/null
+++ b/tests/test_results.py
@@ -0,0 +1,39 @@
+"""Unit tests for ``backend.results`` ? the Success/Failure contract + the
+``unwrap_or_failure``
+helper that standalone scripts use to avoid ``len()``-ing a raw ``Success`` dataclass."""
+import pytest
+
+from backend.results import Failure, Success, unwrap_or_failure
+
+def test_unwrap_success_returns_value_and_no_failure():
+    value, failure = unwrap_or_failure(Success(value=[{"start": 1.0}]))
+    assert value == [{"start": 1.0}]
+    assert failure is None
+
+def test_unwrap_failure_returns_failure_and_no_value():
+    f = Failure(message="Could not open video: x.mp4", is_recoverable=False)
+    value, failure = unwrap_or_failure(f)
+    assert value is None
+    assert failure is f
+
+def test_unwrap_bare_value_passes_through():
+    """Legacy segmenters that predate the Result contract still return a bare list ? the helper
+    passes those through unchanged as ``(value, None)`` so a mixed segmenter set all works."""
+    bare = [{"start": 1.0}, {"start": 2.0}]
+    value, failure = unwrap_or_failure(bare)
+    assert value == bare
+    assert failure is None
+
+def test_unwrap_fixes_the_len_on_success_crash():
+    """Regression for the scripts bug: ``len()`` on a raw ``Success`` raises ``TypeError``;
+    unwrapping first yields the underlying list so callers (run_phase3_eval /
+    run_process_and_stitch, which default to the ``motion`` segmenter) can ``len()`` it."""
+    result = Success(value=[1, 2, 3])
+    with pytest.raises(TypeError):
+        len(result) # the original bug (Success has no __len__)
+    segs, failure = unwrap_or_failure(result)
+    assert failure is None
+    assert len(segs) == 3 # fixed

```

4. assistant (2026-06-18T02:00:09.560Z)

Now let me read the touched files in full context and the related code (especially the inline unwrap in main.py that this DRYs).

5. user (2026-06-18T02:00:10.503Z)

```
1      """Lightweight Result type for standardizing success vs. failure across the pipeline.
2
3      This is the start of the Error/Result Contract standardization (low-regret item #3).
4
5      The goal is to make it impossible to confuse "no data / no rallies" with "something
6      actually failed"
7      (e.g. missing API key, WSL healthcheck failure, provider error, etc.).
8
9      We use simple dataclasses so callers can use isinstance checks or structural pattern
10     matching.
11
12     This is intentionally small and incremental ? we start by applying it to the segmenter
13     and provider layers (the biggest sources of silent [] returns today).
14
15     The types are generic where it makes sense. Callers can still access .value on Success
16     or .message on Failure.
17
18     This module has no dependencies on the rest of the app.
19     """
20
21     from __future__ import annotations
22
23     from dataclasses import dataclass, field
24     from typing import Generic, TypeVar, List, Optional, Any, Tuple, Union
25
26     T = TypeVar("T")
27
28     @dataclass
29     class Success(Generic[T]):
30         """Represents a successful outcome with an optional list of non-fatal warnings."""
31
32         value: T
33         warnings: List[str] = field(default_factory=list)
34
35         def is_success(self) -> bool:
36             return True
37
38         def unwrap(self) -> T:
39             """Return the value. Only safe to call if you already know it's a Success."""
40             return self.value
41
42     @dataclass
43     class Failure:
44         """Represents a failure with a human-readable message and optional details."""
45
46         message: str
47         details: Optional[dict[str, Any]] = None
48         is_recoverable: bool = False # e.g. missing API key (recoverable) vs internal crash
49
50         def is_success(self) -> bool:
51             return False
52
53         def unwrap(self) -> Any:
54             """Raise to make misuse obvious in code that assumes success."""
55             raise RuntimeError(f"Attempted to unwrap a Failure: {self.message}")
56
57     # Convenient alias for use in type hints: Result[list[dict], Failure]
58     # Result[ValueType] means Success[ValueType] | Failure
59     # (the Failure "error" type is fixed for this contract)
60     Result = Union[Success[T], Failure]
61
```

```

62
63     def unwrap_or_failure(result: Any) -> Tuple[Optional[Any], Optional[Failure]]:
64         """Split a segmenter/provider return into ``(value, failure)`` so a caller branches
once.
65
66         - ``Failure``          -> ``(None, failure)``
67         - ``Success(value=v)`` -> ``(v, None)``
68         - a bare value (legacy code that predates this contract and still returns a raw
list)
69         -> ``(value, None)``
70
71         This is the same isinstance/``.value`` unwrap that ``backend/main.py`` does inline;
callers
72         that forget it (e.g. ``len(process_video(...))``) crash with ``TypeError: ... has no
len()``
73         on a ``Success`` dataclass. ``result`` is typed ``Any`` because the contract is
heterogeneous
74         (Success | Failure | legacy bare value)."""
75         if isinstance(result, Failure):
76             return None, result
77         value = result.value if hasattr(result, "value") else result
78         return value, None
79

```

6. assistant (2026-06-18T02:00:13.111Z)

Now let me look at the inline unwrap in main.py that this claims to mirror, and check how the scripts use the result.

7. user (2026-06-18T02:00:15.066Z)

```

35:from backend.results import Failure # standardized error/result contract
192:     result = segmenter.process_video(video_id, req.video_path, req.player_pool,
req.max_frames)
195:     # state (same destroy-after-confirm contract as the Failure branch), then let
201:     if isinstance(result, Failure):
213:     segments = result.value if hasattr(result, "value") else result
220:     "message": f"Successfully indexed {len(segments)} play segments using
'{seg_name}' segmenter.",
347:def process_video(req: ProcessRequest):
835:     result = segmenter.process_video(video_id, args.video_path,
max_frames=args.max_frames)
838:     if isinstance(result, Failure):
843:     segments = result.value if hasattr(result, "value") else result

```

8. assistant (2026-06-18T02:00:17.707Z)

Let me read both inline unwrap sites in main.py to confirm the helper faithfully mirrors them.

9. user (2026-06-18T02:00:18.275Z)

```

188         # if the run yields segments; otherwise drop the new partial rows and keep the
old.
189         play_wm, cand_wm = db_instance.segment_watermarks(video_id)
190         segmenter = segmenter_cls(db_instance, global_config)
191         try:
192             result = segmenter.process_video(video_id, req.video_path, req.player_pool,
req.max_frames)
193         except Exception:
194             # A raising segmenter must not leave half-written rows: restore the pre-run
195             # state (same destroy-after-confirm contract as the Failure branch), then
let
196             # the job worker record the failure.
197             db_instance.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)

```

```

198         raise
199     segmentation_ran = True
200
201     if isinstance(result, Failure):
202         logger.error(f"[API] Segmenter failed: {result.message}")
203         db_instance.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
204         response = {
205             "video_id": video_id,
206             "status": "failed",
207             "message": f"Segmenter '{seg_name}' failed: {result.message}",
208             "results": [r.model_dump() for r in val_results],
209             "play_segments": [],
210         }
211         return response
212
213     segments = result.value if hasattr(result, "value") else result
214     notify("finalizing")
215     _finalize_reingest(video_id, segments, play_wm, cand_wm)
216
217     response = {
218         "video_id": video_id,
219         "status": "success",
220         "message": f"Successfully indexed {len(segments)} play segments using
'{seg_name}' segmenter.",
221         "results": [r.model_dump() for r in val_results],
222         "play_segments": segments,
223     }
224     return response
225 finally:
226     # Always emit the per-video observability report (next to the video, or to
227     # report.output_dir). Never raises. Surface the paths in the response.

```

10. assistant (2026-06-18T02:00:21.036Z)

The helper exactly mirrors the inline pattern at line 201/213: `isinstance(result, Failure)` then `result.value if hasattr(result, "value") else result`. Let me read the second inline site (line 835-843) to confirm it's the same, and check the full context of the two scripts to verify the change is correct and doesn't introduce issues.

11. user (2026-06-18T02:00:21.915Z)

```

830
831     print(f"[CLI] Indexing play segments using segmenter: {seg_name}...")
832     segmenter_actual = seg_name
833     play_wm, cand_wm = db_instance.segment_watermarks(video_id)
834     segmenter = segmenter_cls(db_instance, global_config)
835     result = segmenter.process_video(video_id, args.video_path,
max_frames=args.max_frames)
836     segmentation_ran = True
837
838     if isinstance(result, Failure):
839         print(f"[FAIL] Segmenter '{seg_name}' failed: {result.message}")
840         print("[FAIL] Indexing halted for video.")
841         segments = []
842     else:
843         segments = result.value if hasattr(result, "value") else result
844         print(f"[SUCCESS] Process complete. Indexed {len(segments)} play
segments inside database: {db_path}")
845         # Destroy-AFTER-confirm: keep new on success, restore prior on
empty/failure.
846         _finalize_reingest(video_id, segments, play_wm, cand_wm)
847     finally:
848         # Always emit the per-video observability report. Never raises.
849         paths = emit_indexer_report(

```

12. assistant (2026-06-18T02:00:26.994Z)

Both inline sites match the helper exactly. Now let me read the full context of the two scripts to verify the edits integrate correctly ? particularly the control flow around the changed lines.

13. user (2026-06-18T02:00:27.988Z)

```
130         if not os.path.exists(csv_path):
131             skipped.append((vid, "no GT csv"))
132             print(f"[skip] {vid}: GT csv missing ({csv_path})")
133             continue
134
135         video_id = get_file_md5(video_path)
136
137         # Segment unless score-only or already present (and not --reprocess).
138         if not args.score_only:
139             already = db.query_play_segments(video_id, {}) or
db.query_candidate_segments(video_id)
140             if already and not args.reprocess:
141                 print(f"[have] {vid}: detections already in DB (use --reprocess to
redo)")
142             else:
143                 print(f"[proc] {vid}: segmenting with {args.segmenter}...")
144                 db.add_video(video_id, video_path, "processed", [])
145                 t0 = time.time()
146                 # process_video returns a Result (Success|Failure) for some segmenters
147                 # (e.g. the default `motion`); unwrap before len() so this doesn't crash
with
148                 # TypeError on a Success dataclass ? mirrors the unwrap in
backend/main.py.
149                 segs, failure = unwrap_or_failure(
150                     segmenter.process_video(video_id, video_path,
max_frames=args.max_frames))
151                 if failure is not None:
152                     print(f"        -> FAILED: {failure.message}")
153                     skipped.append((vid, f"segment failed: {failure.message}"))
154                     continue
155                 print(f"        -> {len(segs)} segments in {time.time() - t0:.0f}s")
156
157         gt = load_annotations(csv_path)
158         report = evaluate(db, video_id, gt, stages, iou_threshold=args.iou,
detector=args.detector)
159         print(format_report_text(report, show_failures=args.show_failures))
160         print()
161         d = report_to_dict(report)
162         d["collector_video_id"] = vid
163         per_video.append(d)
164
```

14. assistant (2026-06-18T02:00:36.416Z)

I see something worth examining. In `run\_phase3\_eval.py`, the `segs` variable is only used inside the `else` block (for the `print(f"-> {len(segs)} segments...")`). The actual evaluation at line 158 reads from the DB (`evaluate(db, video\_id, ...)`), not from `segs`. So `segs` is only used for the log line. That's fine ? the unwrap is purely so `len(segs)` doesn't crash.

But let me check: previously when `process\_video` returned a `Failure` (e.g. motion segmenter fails), the old code `segs = segmenter.process\_video(...)` would then crash at `len(segs)` because `Failure` also has no `\_\_len\_\_`. The new code handles failure by `continue`. Good ? but note the `continue` skips evaluation entirely now, whereas before it would crash. That's an improvement. Let me verify the motion segmenter actually returns a Result type to confirm the bug is real.

15. user (2026-06-18T02:00:36.648Z)

```
backend\pipeline\segmenters\yolo_hybrid.py:57: def process_video(self, video_id: str,
video_path: str, player_pool: List[str] = None, max_frames: int = None) -> List[Dict[str, Any]]:
backend\pipeline\segmenters\motion.py:20: def process_video(self, video_id: str, video_path:
str, player_pool: List[str] = None, max_frames: int = None) -> List[Dict[str, Any]]:
backend\pipeline\segmenters\motion.py:31: return Failure(message=f"Could not open
video: {video_path}", is_recoverable=False)
backend\pipeline\segmenters\motion.py:224: return Success(value=segments_data)
backend\pipeline\segmenters\base.py:24: def process_video(self, video_id: str, video_path:
str, player_pool: List[str] = None, max_frames: int = None) -> "Result[List[Dict[str, Any]]]":
backend\pipeline\segmenters\gemini.py:35: def process_video(self, video_id: str, video_path:
str, player_pool: List[str] = None, max_frames: int = None) -> List[Dict[str, Any]]:
backend\pipeline\segmenters\trajectory_hybrid.py:106: def process_video(self, video_id: str,
video_path: str, # type: ignore[override]
```

16. assistant (2026-06-18T02:00:39.848Z)

Confirmed the bug is real: `motion.py` (the default segmenter) returns `Success(value=...)` at line 224 and `Failure(...)` at line 31. The old scripts did `len(segs)` directly on that ? `TypeError`. Let me confirm `motion` is indeed the default segmenter, and check the second script's full flow.

17. user (2026-06-18T02:00:40.715Z)

```
60         print("[CACHE] Using cached play segments. Skipping segmenter API
call.", flush=True)
61         cache_hit = True
62     else:
63         print("[CACHE] Ignoring cache...", flush=True)
64     else:
65         print("[CACHE] Video exists in DB but contains 0 play segments. Treating as
cache miss.", flush=True)
66
67     if not cache_hit:
68         print("[CACHE] Initializing clean indexing...", flush=True)
69         db.clear_video_data(video_id)
70
71         print("Bypassing OpenCV validations for performance on massive files...",
flush=True)
72         db.add_video(video_id, video_path, "processed", [])
73
74         print("Running segmenter & indexing...", flush=True)
75         seg_name = cfg.indexing.default_segmenter
76         segmenter_cls = segmenter_registry.get(seg_name)
77         if not segmenter_cls:
78             print(f"[FAIL] Segmenter '{seg_name}' not found.")
79             return
80         print(f"Using segmenter: {seg_name}", flush=True)
81         segmenter = segmenter_cls(db, cfg)
82         # process_video returns a Result (Success|Failure) for some segmenters (e.g.
`motion`);
83         # unwrap before len()/truthiness so this doesn't crash with TypeError on a
Success
84         # dataclass (and so the empty-guard below actually sees the list) ? mirrors
backend/main.py.
85         segments, failure = unwrap_or_failure(segmenter.process_video(video_id,
video_path))
86         if failure is not None:
87             print(f"[FAIL] Segmentation failed: {failure.message}")
88             return
89         print(f"Successfully indexed {len(segments)} play segments.", flush=True)
90
91     if not segments:
92         print("No play segments detected. Cannot compile highlights.")
```

```

93         return
94
95
96     end_padding = cfg.stitching.end_padding
97     begin_padding = cfg.stitching.begin_padding
98     min_shot_count = cfg.stitching.min_shot_count
99

```

18. assistant (2026-06-18T02:00:44.829Z)

Now there's a subtle thing to check here. In `run\_process\_and\_stitch.py`, `segments` is referenced at line 91 (`if not segments:`) regardless of whether `cache\_hit` was true. Let me check what `segments` is in the cache-hit path? if `cache\_hit` is true, the `if not cache\_hit:` block at line 67 is skipped, so `segments` must be assigned elsewhere. Let me read the start of the function.

19. user (2026-06-18T02:00:45.825Z)

```

30     cfg = load_config() # typed AppConfig ? single source of defaults
31     # Re-apply now that the verbose_http knob is known (default keeps HTTP chatter
quiet).
32     setup_logging(verbose_http=cfg.logging.verbose_http)
33
34     print("Verifying files exist...")
35     if not os.path.exists(video_path):
36         print(f"Error: Video file not found at {video_path}")
37         return
38
39     os.makedirs(cfg.output.output_dir, exist_ok=True)
40     db_path = os.path.join(cfg.output.output_dir, cfg.output.db_name)
41     db = Database(db_path)
42
43     # Calculate MD5 of the video to uniquely identify it
44     import time
45     print("Calculating video MD5 (this might take a few seconds on large files)...",
flush=True)
46     t0 = time.time()
47     video_id = compute_video_id(video_path)
48     print(f"Calculated MD5: {video_id} (took {time.time()-t0:.1f}s)", flush=True)
49
50     segments = []
51     cached_video = db.get_video(video_id)
52     cache_hit = False
53
54     if cached_video and cached_video.get("status") == "processed":
55         segments = db.query_play_segments(video_id, {})
56         if len(segments) > 0:
57             print(f"\n[CACHE] Found cached indexed video in database (MD5: {video_id})
with {len(segments)} parsed play segments.", flush=True)
58             choice = input("Would you like to (1) Run stitching on the cached segments,
or (2) Reprocess the video with Gemini from scratch? [1/2]: ").strip()
59             if choice == '1':
60                 print("[CACHE] Using cached play segments. Skipping segmenter API
call.", flush=True)
61                 cache_hit = True
62             else:
63                 print("[CACHE] Ignoring cache...", flush=True)
64         else:

```

20. assistant (2026-06-18T02:00:51.627Z)

`segments` is initialized to `[]` at line 50 and assigned from the DB query in the cache-hit path (line 55), so the `if not segments:` guard at line 91 is safe in all branches. The change is correct.

Now let me note one important behavioral pre-existing detail: in the cache-hit path, `segments` comes from `db.query\_play\_segments` (a raw list) ? that path never went through `process\_video`, so it was never affected by the bug, and the PR correctly doesn't touch it.

Let me verify the test file imports and that the tests align with the actual behavior, then run the new test plus the full suite to satisfy the pre-LGTM gate.

21. user (2026-06-18T02:00:52.114Z)

```
1      """Unit tests for ``backend.results`` ? the Success/Failure contract + the
  ``unwrap_or_failure``
2      helper that standalone scripts use to avoid ``len()``-ing a raw ``Success``
  dataclass."""
3      import pytest
4
5      from backend.results import Failure, Success, unwrap_or_failure
6
7
8      def test_unwrap_success_returns_value_and_no_failure():
9          value, failure = unwrap_or_failure(Success(value=[{"start": 1.0}]))
10         assert value == [{"start": 1.0}]
11         assert failure is None
12
13
14     def test_unwrap_failure_returns_failure_and_no_value():
15         f = Failure(message="Could not open video: x.mp4", is_recoverable=False)
16         value, failure = unwrap_or_failure(f)
17         assert value is None
18         assert failure is f
19
20
21     def test_unwrap_bare_value_passes_through():
22         """Legacy segmenters that predate the Result contract still return a bare list ? the
  helper
23         passes those through unchanged as ``(value, None)`` so a mixed segmenter set all
  works."""
24         bare = [{"start": 1.0}, {"start": 2.0}]
25         value, failure = unwrap_or_failure(bare)
26         assert value == bare
27         assert failure is None
28
29
30     def test_unwrap_fixes_the_len_on_success_crash():
31         """Regression for the scripts bug: ``len()`` on a raw ``Success`` raises
  ``TypeError``;
32         unwrapping first yields the underlying list so callers (run_phase3_eval /
33         run_process_and_stitch, which default to the ``motion`` segmenter) can ``len()``
  it."""
34         result = Success(value=[1, 2, 3])
35         with pytest.raises(TypeError):
36             len(result) # the original bug (Success has no __len__)
37         segs, failure = unwrap_or_failure(result)
38         assert failure is None
39         assert len(segs) == 3 # fixed
40
```

22. assistant (2026-06-18T02:00:59.529Z)

Let me consider an edge case in the helper: a `Success` with `value=None`. The helper returns `(None, None)`. A caller checking `if failure is not None` would proceed with `value=None`, and then `len(None)` would crash. But `Success.value` for segmenters is always a list, so this is theoretical. Still, worth noting `value=None` and `Failure` both could lead to `value is None` ? but the helper distinguishes via the `failure` slot, so callers branch on `failure is not None`, which is correct. No issue.

Now let me check one more subtle thing ? the bare-value `hasattr(result, "value")` path. A plain list has no `.value` attribute, so it passes through. But what about a dict that happens to have a `"value"` key? `hasattr` checks attributes, not dict keys, so a dict passes through correctly. Good.

Let me run the new test and the full suite to satisfy the verification gate.

23. user (2026-06-18T02:01:01.761Z)

```
.... [100%]
4 passed in 0.07s
```

24. assistant (2026-06-18T02:01:04.525Z)

New tests pass. Now let me run the full suite to confirm no regression, and lint/typecheck the touched files.

25. user (2026-06-18T02:02:24.104Z)

```
=====
Full test suite: C:\Users\avidu\AppData\Local\Programs\Python\Python313\python.exe -m pytest
-q
obsreport (optional reporting dep): available
=====
.....SS..... [ 8%]
..... [ 16%]
..... [ 25%]
..... [ 33%]
..... [ 42%]
..... [ 50%]
..... [ 58%]
..... [ 67%]
..... [ 75%]
.SSSSS..... [ 84%]
..... [ 92%]
..... [100%]
850 passed, 7 skipped in 74.70s (0:01:14)
=====
RESULT: PASS - all tests green
=====
```

26. assistant (2026-06-18T02:02:26.067Z)

Full suite green (850 passed, 7 skipped). Now lint and typecheck the touched files to match CI.

27. user (2026-06-18T02:02:28.746Z)

```
All checks passed!
---MYPY---
Success: no issues found in 1 source file
```

28. assistant (2026-06-18T02:02:36.772Z)

Lint and typecheck clean. Let me check whether the repo has a coverage floor in CI that I should verify against, and confirm `motion` is the default segmenter (the PR claims scripts "default to motion", crashing).

29. user (2026-06-18T02:02:38.958Z)

```
./.github/workflows/ci.yml:97: run: python run_all_tests.py --cov=backend --cov=training
--cov-report=term --cov-fail-under=70
---DEFAULT SEGMENTER---
backend/config/models.py:118: default_segmenter: str = "yolo_hybrid"
backend/config/models.py:192: @field_validator("default_segmenter")
```

Binary file backend/config/__pycache__/models.cpython-313.pyc matches

30. assistant (2026-06-18T02:02:42.340Z)

Interesting finding to verify. The default segmenter in config is `yolo\_hybrid`, NOT `motion`. Let me check whether `yolo\_hybrid` returns a Result type or a bare list, because the PR's framing and the script comments say "default to motion". This affects whether the crash actually happens under the default config.

31. user (2026-06-18T02:02:43.177Z)

```
57         def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None, max_frames: int = None) -> List[Dict[str, Any]]:
58             self.person.lazy_load_model()
59
60             # Get sport profile config
61             sport_name = self.config.get("sport", "badminton")
62             try:
63                 sport_profile = sport_registry.get(sport_name)
64             except KeyError as e:
65                 logger.error(str(e))
66                 return []
67
68             # Get AI provider and model config
69             idx_cfg = self.config.get("indexing", {})
70             provider_name = self.config.get("ai_provider", idx_cfg.get("ai_provider",
"gemini"))
71             model_name = self.config.get("ai_model", idx_cfg.get("ai_model",
"gemini-2.5-flash"))
72
73             # Get appropriate API key based on the provider
74             api_key_env_var = f"{provider_name.upper()}_API_KEY"
75             api_key = os.environ.get(api_key_env_var)
76             if not api_key:
77                 logger.error(
78                     f"{api_key_env_var} environment variable is not set! "
79                     f"To run the {provider_name} segmenter, set your API key. "
80                     f"In PowerShell: $env:{api_key_env_var}=\"your_api_key_here\""
81                 )
82                 return []
83
84             try:
85                 provider = provider_registry.get(provider_name, api_key=api_key,
model_name=model_name)
86             except KeyError as e:
87                 logger.error(str(e))
88                 return []
89
90             logger.info(f"Initializing Hybrid Segmenter (torchvision + ByteTrack) for video:
{video_path}")
91
92             video_duration = get_video_duration(video_path)
93             if video_duration <= 0:
94                 logger.error("Invalid video duration.")
95                 return []
96
```

32. assistant (2026-06-18T02:02:48.176Z)

`yolo\_hybrid` (the actual config default) returns bare lists, not a Result. So under the *config* default, the old scripts would NOT crash. The crash happens under the `motion` segmenter. Let me check what the scripts actually default to ? `run\_phase3\_eval.py` takes a `--segmenter` arg, and `run\_process\_and\_stitch.py` uses `cfg.indexing.default\_segmenter`. Let me check the `--segmenter` default in run_phase3_eval.

33. user (2026-06-18T02:02:48.888Z)

```
12:         --segmenter motion --only shuttleset_1
15:     python run_phase3_eval.py --manifest ../manifest.json --segmenter yolo_hybrid \
51:     p.add_argument("--videos-root", default=None,
54:     p.add_argument("--segmenter", default="motion",
57:     p.add_argument("--stage", choices=("candidates", "final", "both", "auto"),
default="auto",
59:     p.add_argument("--detector", default=None, help="Restrict candidate scoring to this
detector.")
60:     p.add_argument("--iou", type=float, default=0.5, help="IoU match threshold (default:
0.5).")
61:     p.add_argument("--only", action="append", default=None,
63:     p.add_argument("--limit", type=int, default=None, help="Process at most N videos.")
64:     p.add_argument("--max-frames", type=int, default=None,
70:     p.add_argument("--db", default=None, help="Indexer DB path (default: output/<db_name>).")
71:     p.add_argument("--json", dest="json_out", default=None, help="Write a JSON report here.")
```

34. assistant (2026-06-18T02:02:56.620Z)

This clarifies it. In `run_phase3_eval.py`, `--segmenter` defaults to `"motion"` (line 54), so running that script with no `--segmenter` flag hits the bug. That's a real, default-path crash for *that script*.

For `run_process_and_stitch.py`, it uses `cfg.indexing.default_segmenter` which is `yolo_hybrid` (bare list) under default config ? so that script would NOT crash under default config. But the bug is still latent: if a user sets `default_segmenter: motion` (or any future Result-returning segmenter), it crashes. The fix is correct and forward-safe regardless.

The script-comment phrasing "for some segmenters (e.g. `motion`)" in `run_process_and_stitch.py` is accurate (it says "some segmenters", not "the default"). The `run_phase3_eval.py` comment is also accurate. The test docstring at line 32-33 says both scripts "default to the `motion` segmenter" ? that's accurate for `run_phase3_eval.py` (its own `--segmenter` default is `motion`) but technically imprecise for `run_process_and_stitch.py` (whose default comes from config = `yolo_hybrid`). This is a test-docstring nit, not a code defect.

Let me double-check there isn't another caller in the codebase that does the same raw-`len` pattern on `process_video` and got missed (scope completeness).

35. user (2026-06-18T02:02:58.163Z)

```
tests\test_api_endpoints.py:55:     def process_video(self, video_id, video_path, *a, **k):
scripts\run_process_and_stitch.py:85:         segments, failure =
unwrap_or_failure(segmenter.process_video(video_id, video_path))
tools\generate_highlights.py:167:         result = seg_cls(db, cast(Dict[str, Any],
cfg)).process_video(vid, proc_detect)
tests\test_async_jobs.py:63:     def process_video(self, video_id, video_path, *a, **k):
tests\test_async_jobs.py:72:     def process_video(self, video_id, video_path, *a, **k):
tests\test_async_jobs.py:170:     def process_video(self, video_id, video_path, *a, **k):
tests\test_base.py:15:     def process_video(self, video_id, video_path, player_pool=None,
max_frames=None) -> Result:
tests\test_base.py:34:     success = inst.process_video("v1", "path")
tests\test_base.py:40:     failure = inst.process_video("v1", "fail")
tests\test_base.py:82:     def process_video(self, video_id, video_path, player_pool=None,
max_frames=None) -> Result:
tests\test_base.py:108:     def process_video(self, video_id, video_path, player_pool=None,
max_frames=None) -> Result:
tests\test_base.py:187:     def process_video(self, video_id, video_path, player_pool=None,
max_frames=None) -> Result:
tests\test_base.py:192:     inst.process_video("v1", "path")
scripts\run_phase3_eval.py:150:         segmenter.process_video(video_id, video_path,
max_frames=args.max_frames))
backend\main.py:192:         result = segmenter.process_video(video_id, req.video_path,
req.player_pool, req.max_frames)
backend\main.py:347: def process_video(req: ProcessRequest):
```

```

backend\main.py:835:         result = segmenter.process_video(video_id, args.video_path,
max_frames=args.max_frames)
tests\test_fusion_hybrid.py:55:         res = seg.process_video("v1", "clip.mp4")
tests\test_gemini.py:42:         res = segmenter.process_video("v1", "dummy.mp4")
tests\test_gemini.py:81:         res = segmenter.process_video("v1", "dummy.mp4")
tests\test_gemini.py:108:         return segmenter.process_video(video_id, "dummy.mp4")
tests\test_gemini.py:185:         res = segmenter.process_video("vid_single_oversize",
"dummy.mp4")
backend\pipeline\segmenters\gemini.py:35:     def process_video(self, video_id: str, video_path:
str, player_pool: List[str] = None, max_frames: int = None) -> List[Dict[str, Any]]:
backend\pipeline\segmenters\motion.py:20:     def process_video(self, video_id: str, video_path:
str, player_pool: List[str] = None, max_frames: int = None) -> List[Dict[str, Any]]:
backend\pipeline\segmenters\base.py:24:     def process_video(self, video_id: str, video_path:
str, player_pool: List[str] = None, max_frames: int = None) -> "Result[List[Dict[str, Any]]]":
backend\pipeline\segmenters\yolo_hybrid.py:57:     def process_video(self, video_id: str,
video_path: str, player_pool: List[str] = None, max_frames: int = None) -> List[Dict[str, Any]]:
backend\pipeline\segmenters\trajectory_hybrid.py:106:     def process_video(self, video_id: str,
video_path: str, # type: ignore[override]
backend\results.py:72:         that forget it (e.g. ``len(process_video(...))``) crash with
``TypeError: ... has no len()``
tests\test_job_worker_loop.py:70:     def process_video(self, video_id, video_path, *a, **k):
tests\test_job_worker_loop.py:135:         def process_video(self, video_id, video_path, *a,
**k):
tests\test_job_worker_loop.py:247:         def process_video(self, *a, **k):
tests\test_tracknet_hybrid.py:63:         res = seg.process_video("v1", "clip.mp4")
tests\test_tracknet_hybrid.py:82:         assert seg.process_video("v1", "clip.mp4") == []
tests\test_tracknet_hybrid.py:95:         res = seg.process_video("v1", "clip.mp4")
tests\test_tracknet_hybrid.py:110:         assert seg.process_video("v1", "clip.mp4") == []
tests\test_tracknet_hybrid.py:129:         seg.process_video("v1", "clip.mp4")
tests\test_tracknet_hybrid.py:153:         res = seg.process_video("v1", "clip.mp4")
tests\test_tracknet_hybrid.py:181:         res = seg.process_video("v1", "clip.mp4")
tests\test_trajectory_hybrid_equivalence.py:56:         return db, seg.process_video("v1",
"clip.mp4")
tests\test_yolo_hybrid.py:89:         res = segmenter.process_video("v1", "dummy.mp4")
tests\test_yolo_hybrid.py:113:         res = segmenter.process_video("v1", "dummy.mp4")
tests\test_yolo_hybrid.py:156:         res = segmenter.process_video("v1", "dummy.mp4")
tests\test_yolo_hybrid.py:225:         res = segmenter.process_video("v1", "dummy.mp4")
tests\test_yolo_hybrid.py:276:         segmenter.process_video("v1", "dummy.mp4")

```

36. assistant (2026-06-18T02:03:02.526Z)

There's one more production caller worth checking: `tools/generate\_highlights.py:167`. Let me check whether it handles the Result type properly or has the same latent bug (it's out of this PR's stated scope, but if it has the identical bug, that's a scope-completeness observation, not a blocker for this PR).

37. user (2026-06-18T02:03:03.788Z)

```

160         if detect_src != src:
161             logger.info("detecting on proxy '%s' -> stitching reel from '%s'",
162                 os.path.basename(detect_src), name)
163             logger.info("video_id=%s segmenting (%s)...", vid, segmenter)
164
165             # AppConfig is consumed by segmenters via its dict-compatible .get shim (same as
166             # main.py); cast keeps the type checker honest about that contract.
167             result = seg_cls(db, cast(Dict[str, Any], cfg)).process_video(vid,
proc_detect)
168             if isinstance(result, Failure):
169                 logger.error("segmentation FAILED: %s", getattr(result, "message",
result))
170                 summary["failed"].append(name)
171                 continue
172
173             segs = db.query_play_segments(vid, {})
174             pairs = sorted((float(s["start_time"]), float(s["end_time"])) for s in segs)

```

```

175         logger.info("rallies detected: %d", len(pairs))
176     if not pairs:
177         logger.warning("no rallies -> no reel")
178         summary["no_rallies"].append(name)
179         continue
180
181     # Stitch source: the original `src` when detecting on a separate proxy (hi-
res reel);
182     # otherwise the (possibly localized) detection input ? identical to legacy
behaviour.
183     stitch_src = src if detect_from_paths else proc_detect
184     # Honor the configured branding watermark + segment paddings (stitching.*).
Read
185     # dict-compatibly so a plain-dict config (tests) and the typed AppConfig
both work.
186     # NOTE: the runner previously dropped these ? reels shipped UN-watermarked
with the
187     # ctor default paddings; this threads config.json through (the on-by-default
mark).
188     st = cfg.get("stitching") or {}
189     VideoCompiler(

```

38. assistant (2026-06-18T02:03:08.084Z)

`tools/generate\_highlights.py` already handles `Failure` correctly (line 168) and reads segments back from the DB rather than `len()`-ing the result, so it does not have the bug. It's a third inline `isinstance(result, Failure)` site that this helper could also DRY, but the PR deliberately scoped to the two crashing scripts ? that's a reasonable scope decision, not a defect. I'll flag it as an optional follow-up nit only.

Let me confirm the commit/PR claims by checking the commit log on this branch.

39. user (2026-06-18T02:03:09.824Z)

03d11d9 fix(scripts): unwrap the segmenter Result before len()/truthiness (motion crashes the default path)

---FILES---

```

backend/results.py
scripts/run_phase3_eval.py
scripts/run_process_and_stitch.py
tests/test_results.py

```

40. assistant (2026-06-18T02:03:19.959Z)

Clean single-commit branch, exactly 4 files touched ? all matching the PR description (the two scripts, the helper in `backend/results.py`, the new test). No stray files, no unrelated changes, no licensing/product/behavior changes.

Let me reconcile one nuance with the commit message: "motion crashes the default path." For `run\_phase3\_eval.py` this is literally true (its `--segmenter` default is `motion`). For `run\_process\_and\_stitch.py`, the default path uses config's `yolo\_hybrid` (bare list ? no crash); motion is opt-in there. So "default path" is accurate for one of the two scripts. Minor imprecision in the commit subject, not a code issue.

I've completed verification. Summary of what I checked:

- The bug is real: `motion.py` returns `Success`/`Failure` (no `\_\_len\_\_`); the old scripts called `len()`/truthiness directly on it.
- The helper faithfully mirrors both inline unwrap sites in `main.py` (lines 201/213 and 838/843): identical `isinstance(Failure)` + `result.value if hasattr(...) else result`.
- Control flow in both scripts is correct after the edit (the cache-hit/`segments=[]` init path in `run\_process\_and\_stitch.py` is unaffected and safe).
- Full suite green (850 passed, 7 skipped), new tests pass, `ruff` and `mypy` clean.
- Scope is exactly right; no other caller has the same latent crash (`generate\_highlights.py` already handles `Failure` and reads from DB).

- Edge cases (``Success(value=None)``, dict with a ``"value"`` key, bare list) all behave correctly.

The only findings are minor accuracy nits in prose (test docstring + commit subject), no code defects.

41. user (2026-06-18T02:03:34.594Z)

Structured output provided successfully

42. assistant (2026-06-18T02:03:38.753Z)

Review complete. PR #213 is a clean, correct bug fix ? I verified the crash is real (motion segmenter returns Success/Failure dataclasses with no ``__len__``), the new ``unwrap_or_failure`` helper faithfully mirrors the existing inline unwrap in ``backend/main.py``, both script edits integrate correctly, and the full suite (850 passed / 7 skipped), new tests, ruff, and mypy are all green. Scope is exactly right (4 expected files, single clean commit). No must-fix findings; the three findings are minor prose-accuracy nits. `blocks_merge=false`.