

claude-windows / 8d4866bf-8edd-4da9-909e-022ddee12675

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\8d4866bf-8edd-4da9-909e-022ddee12675\subagents\workflows\wf_4e32c8b8-725\agent-a2cebae992e871570.jsonl`  
- SHA-256: `00c011930c520d96f321a594686428cb8e1ef78a74e6db192fee8948e0f28103`  
- Source modified: `2026-06-19T17:19:34+00:00`  
- Imported at: `2026-07-05T16:48:25+00:00`  
- project: `wf_4e32c8b8-725`  
- session_id: `8d4866bf-8edd-4da9-909e-022ddee12675`
```

Transcript

1. user (2026-06-19T17:18:10.918Z)

You are auditing the docs/ folder of the badminton-highlight-indexer repo to decide which docs can move to docs/archives/. Be rigorous and EVIDENCE-BASED ? the owner has run two doc-staleness sweeps and cares deeply about not archiving live work or breaking links.

REPO ARCHIVE POLICY (docs/archives/README.md) ? APPLY STRICTLY:

- Archive ONLY work in a clear TERMINAL state: DONE / COMPLETED / REJECTED / fully SUPERSEDED. "Never archive in-flight or pending work."
- KEEP-LIVE (do NOT archive) anything that is: living/active, in-progress, a not-yet-started but still-intended design/plan (owner-gated future work is LIVE, not terminal), a reference/navigation doc (CODE_MAP, README, HOW_RALLY_DETECTION_WORKS, EVALUATION), a template (PROJECT_DOC_TEMPLATE), or a roadmap/backlog/strategy doc still being tracked.
- Destinations inside archives/: past_projects/ (bounded completed project / tech-debt / starter effort) · research/ (completed or superseded research notes/plans) · decisions/ (ADRs) · checkpoints/ (dated state snapshots) · archives/ top-level (one-off audits/reports/launch records). Naming convention for a completed doc: <NAME>-COMPLETED-<date>.md.
- ALREADY-ARCHIVED COPIES EXIST for these (so the top-level file is likely a pointer-stub or a stale DUPLICATE to reduce, NOT a fresh move): docs/archives/CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md, docs/archives/DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md, docs/archives/HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md. If your doc already has an archived copy, check whether the top-level is a stub (leave it) or a full duplicate (flag DEDUPE).
- CONTRADICTION TRAP: a doc may say "COMPLETED" yet still be cited as a LIVE source elsewhere (e.g. CLAUDE.md cites DEFERRED_HARDENING_PLAN.md #16 as the next platform slice). Before recommending ARCHIVE, grep CLAUDE.md, docs/NEXT_STEPS.md, docs/README.md, and other top-level docs for the filename. If it is referenced as a live/next source, verdict = ASK_OWNER (note the contradiction), not ARCHIVE_NOW.

YOUR BATCH (read each from docs/<name>): FIELD_VISIT_ANSWER_SHEET.md, FIELD_VISIT_PLAYBOOK.md, PRODUCT_FIELD_ASSOCIATE_FLYER.md, PMF_FIELD_SIGNALS.md, PMF_MARKET_RESEARCH.md, KHELSUTRA_PRODUCT_OVERVIEW.md

For EACH doc:

1. Read the status header + TL;DR/intro + (if present) the \$7/\$8 progress tracker + definition-of-done. You do NOT need to read every line of long docs ? sample enough to judge lifecycle state honestly.
2. Decide lifecycle and verdict per the strict policy. Remember: owner-gated / not-yet-started DESIGN docs are LIVE future work (KEEP_LIVE), NOT terminal. Only DONE/COMPLETED/REJECTED/fully-superseded is terminal.
3. Check for an already-archived copy (the policy lists three; also Glob docs/archives for the doc's stem). If the top-level is a thin pointer-stub ? KEEP_LIVE (it's intentional). If it's a FULL DUPLICATE of an archived COMPLETED copy ? DEDUPE_STUB (recommend reducing top-level to a stub or removing).
4. Cross-check "cited as live": Grep CLAUDE.md, docs/NEXT_STEPS.md, docs/README.md, and the other top-level docs/ for this filename. If it's cited as a live/next/authoritative source, it

CANNOT be archived even if it says COMPLETED ? ASK_OWNER with the contradiction noted.

5. Inbound references: Grep the whole repo (docs/, backend/, CLAUDE.md, *.py) for the filename to count link-breakage risk; list up to 6 key referrers.

Return the structured object with one entry per doc in your batch. Be decisive: most reference/template/active/design docs are KEEP_LIVE; reserve ARCHIVE_NOW for genuinely terminal single-efforts/reports/completed projects, and ASK_OWNER for the contradictory/ambiguous ones.

2. assistant (2026-06-19T17:18:13.252Z)

I'll audit each doc in my batch. Let me start by reading the key reference files and the docs themselves.

3. user (2026-06-19T17:18:18.474Z)

Your configuration specifies to merge with the ref 'refs/heads/docs/court-floor-rally-end' from the remote, but no such ref was fetched.

---ARCHIVES---

CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md

COMMERCIAL_READINESS_REVIEW.md

DEFERRED_HARDENING_PLAN-COMPLETED-2026-06-14.md

HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md

MONETIZATION_AUDIT.md

README.md

TEST_RUN_SUMMARY.md

checkpoints

decisions

past_projects

quality-iterations

research

---ARCHIVE STEMS for my batch---

4. user (2026-06-19T17:18:18.483Z)

```
1      # Next steps (live) ? owned-model + multisport
2
3      **Updated:** 2026-06-13 (**rally-detection quality track**: fusion + experiment-harness
design docs merged,
4      owner-gated ? see the dedicated section below, the desktop-agent pick-up). Prior:
2026-06-10 (audit?M0?Gemini-battery session; later same day the **UI-alpha track opened**:
5      ADR-009 ratified ? KhelSutra Guru / khelsutra.guru, private `khelsutra` repo, local-
first delivery ? see
6      `docs/UI_ALPHA_EXECUTION_PLAN.md`; **PR-1 async jobs (#16) MERGED 2026-06-11, PR #87**;
**PR-2
7      upload/download MERGED 2026-06-11, PR #99** ? live-E2E-verified (3-part chunked upload
MD5 byte-identity +
8      compile?browser-download); next: PR-3 identity. Same day: the **8?9 quality sweep
#90?#98 landed** ? ruff/mypy/
9      coverage-floor CI lanes, SQLite conn-close fix, config unknown-key warning, hybrid-twin
collapse into
10     `trajectory_hybrid.py`, WSL tilde+abspath live-path fixes; suite 388 green; LOVO 0.626
re-verified exact).
11     **Authoritative roadmap:**
12     `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` (+ ADR-008 for topology). **Latest blow-by-
blow:** memory
13     `current-state.md` / `session-handoff.md`. This file = the prioritized "what to do
next," refreshed each session.
14
15     ## ? Next-session / GPU-box pickup (2026-06-15)
16     **M0 in-container quick wins are COMPLETE:** ship-gate provisional blocker (**#170**),
train/eval split guard (**#171**), the **doc-consistency sweep** (**#174** ? corrected 16
drifted docs vs the code), and the **P3 person-feature fusion litmus machinery is already
present + suite-green** (`backend/eval/fusion_compare.py` + `ablation.py`: person-driven ?F1 /
bootstrap CI / paired sign test; the real-golden `skipif` litmus waits only on GPU-extracted
```

data). ? CI is red repo-wide due to a ****GitHub Actions minutes/billing outage**** (account-level, not code) ? work is local-gated + `--admin`-merged` until Actions is restored (Settings ? Billing ? Actions minutes / spending limit).

17

18 ****Immediate next step ? the real P3 LOVO measurement**** (this box has a GPU + ``torch+cu124` + `cv2``; videos and trajectory CSVs are present, but the GPU-extracted golden features + ``.rallies.csv`` labels are not yet here):

19 1. ****[owner / GPU box]**** ``python -m backend.eval.fusion_golden --manifest output/human_lovo_manifest.json --out-dir "%USERPROFILE%\OneDrive\rally-annotator\golden-data\features\2026-06-15-n6`` ? produces ``.golden_features.json`` for the 6 golden videos. (GPU + WSL + videos + trajectory CSVs + ``.rallies.csv`` labels are owner-side only.)

20 2. Artifacts land in the ****shared cloud folder**** ``C:\Users\avidu\OneDrive\rally-annotator\golden-data\`` (OneDrive auto-syncs both ways). Convention + workflow: ****[`GOLDEN\_DATA\_SHARING.md`](GOLDEN_DATA_SHARING.md)****; tracked in ****#175****.

21 3. ****[CPU dev/agent session]**** ``python -m backend.eval.fusion_compare --features-dir <shared>\features\2026-06-15-n6 --record`` + ``python -m backend.eval.ablation --features-dir <shared>\... --granularity group`` ? real person-feature ?F1; copy the JSONs into ``.output/`` to also pass the ``.skipif`` litmus tests (``.tests/test_ablation.py::test_litmus_reproduces_gate_a``).

22

23 ****Also queued (owner-gated / pending input):****

24 - Ship-gate target calibration ? ****#172**** (owner decision; needs corpus growth).

25 - Tests reorganization ? blocked on the owner's grouping metric (provide it ? an agent will do it).

26 - Restore CI ? the Actions outage is account-level (billing/minutes), not a repo/config bug.

27

28 ****Leaving (bulky / M2 ? no start without owner go):**** ActionFormer (M0.4), event-index DB migration (M0.1), Gemini-silver purge (M0.3), cost-to-Gen-2 benchmark + ByteTrack deprecation, multisport, repo rename, PMF execution.

29

30 **## Where we are (one paragraph)**

31 ADR-008 ratified ("repo split driven by productionization, not isolation"): training lives in-repo (``.training/``)

32 behind a CI-enforced import wall; the featurizer is single-sourced (``.backend/features/rally_seq.py``, `train==serve`);

33 the ****Gen-0 harness ships**** (``python -m training.gen0.harness --manifest output/human_lovo_manifest.json --floor eval_baselines/heuristic_lovo_n6.json --version <semver>``) and produced artifact v0.1.0 (LOVO ****0.626****, floor passed,

35 sign test 5/6 wins $p=0.109$? honestly not significant, `ship=False`). The ****n=6** owned corpus is in the collector

36 (262 rallies, Proprietary, commercial export = owned data only after the ShuttleSet reclassification). The ****Gemini**

37 battery is done** (see table) ? the moat is quantified.

38

39 **## The quality table that now drives strategy (IoU 0.5 vs human golden labels)**

40 | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |

41 |---|---|---|

42 | Heuristic (velocity windowing) | 0.657 | 0.157 |

43 | ****Gen-0 owned head**** (LOVO, n=6) | 0.744 | 0.561 |

44 | Two-stage WASB+Gemini refine | 0.83 | ? |

45 | ****Gemini wrapper**** (``.gemini-flash-latest``) | ****0.93**** | ****0.66**** |

46

47 ****Reading:**** clean footage is commoditized (serve it with Gemini for cents). ****Multi-court** drops the commodity

48 wrapper to 0.66 ? that is the owned model's ship target** (its FPs are background-game pickups + dead-time merges,

49 the exact contrast-metric confound). Gen-0 is 0.10 behind with only n=6 ? a corpus-growth-sized gap, not an

50 architecture-sized one. Two-stage refine = the cost-efficient serving path for LONG tapes (uploads candidate clips

51 only); wrapper 503-flakiness (observed all session) makes the provider-fallback registry load-bearing.

52 Baselines tracked in ``.eval_baselines/`` (heuristic_lovo_n6.json, gemini_wrapper_2026-06-10.json).

53

54 ## Rally-detection quality track (design merged 2026-06-13, owner-gated; implementation in progress)

55 **This is the track a desktop/local agent is slated to pick up.** Design docs landed (PR #112 + session). All implementation is **owner-gated** (per explicit in the plan); each step = ONE PR off `master`.

56

57 **Hardening loop (T1-T5 positive tests + audit + finale) COMPLETE (2026-06-14, PRs #156/#157/#160/#161/#163 + #165 + final records).** See archived `docs/archives/CODE\_AUDIT\_AND\_TEST\_HARDENING-COMPLETED-2026-06-14.md` (Living Log with every pre/post/review) and `docs/archives/HARDENING\_LOOP\_HANDOFF-COMPLETED-2026-06-14.md` (full rules + STATE). Top-level now pointer stubs per #165 review. All 5 items + verification + archive done; clean slate.

58

59 - **Docs:** `docs/MULTI\_SIGNAL\_FUSION\_PLAN.md` (shuttle + person + audio, one fusion layer) .

60 `docs/EXPERIMENT\_HARNESS.md` (A/B + shadow + promotion gate, zero-regression) .

61 `docs/ROORA\_LABELING\_GUIDELINES.md` (v8) . `docs/HOW\_RALLY\_DETECTION\_WORKS.md` (2-layer mental model).

62 - **Key finding (no re-investigation needed):** the owner's "held-shuttle counted as a rally" bug is **NOT** a missing capability ? `backend/pipeline/detectors/rally\_gate.py` already implements the net-crossing **"Gate A"** (`apply\_gate(..., min\_crossings=2)`) that kills service-feed/held-shuttle windows, but it is **only** called from eval drivers (and there was no `min\_crossings` knob in `IndexingConfig`). **Wiring Gate A into the live segmenter (now via fusion_hybrid)** was the single cheapest, highest-confidence quality win. (P2 FusionSegmenter specifics: registry `fusion\_hybrid`, **default-OFF**, seeds from substrate, persists `net\_crossings` + optional person features, optional served Gate A/B via config knobs; adversarially reviewed; suite green. See `tests/test\_served\_gate\_a\_regression.py` for the honest finding that on production windowing Gate A is neutral/negative in some cases, hence kept opt-in with knob for safety, and the leaderboard revert note.)

63 - **Current status (2026-06-14):** P1 (person cue extraction to reusable `PersonDetector`) + P2 (FusionSegmenter + Gate A signal + served handoff gating) + experiment harness P1 DONE (adversarial review + NO-REGRESSION; suite green). **P3 next** (this PR + follow-ups): learned fusion (person features ? harness `compare` LOVO lift on the 6 golden videos; eager-person-compute optimisation); then P4 audio via hit_detector.

64 - **Pick-up order** (each = ONE PR off `master`, owner approval first, flag-gated default-OFF, promote only on measured LOVO):

65 1. (P2 first-step, largely done via fusion_hybrid) Wire Gate A into production served path + config knob (min_crossings); apply in hybrids at runtime. (Now available opt-in; served neutral on production windowing per litmus ? kept opt-in with knob for safety.)

66 2. P3 first step (this work item): harness `compare` LOVO litmus for person-on variant (players_in_court + two_sided + colocation etc. from `fusion\_features`) on the 6 golden; add the eager compute optimisation note/landing. Pure + harness (testable here).

67 3. P3/P4 continuation: audio cue integration + full learned promotion if lift proven.

68 - **Discipline (from EXPERIMENT_HARNESS.md):** every new cue **flag-gated**, default **OFF**; promote to default **only** on measured LOVO lift through `run\_regression.py` (exit 0/1/3); pinned `CONTROL` ? rollback = config flip. Harness re-uses ~80% existing code.

69 - ? Gate A, person cue logic, harness, and pure fusion paths are testable here; the heavier real-video / golden LOVO / ablation confirmation waits on the GPU-extracted golden features + `rallies.csv` labels, which are not yet present in this checkout.

70 - **External due-diligence review inputs (2026-06-17) ? PICK UP AFTER** current ideas are exhausted; full details + citations in [`RALLY\_QUALITY\_RESEARCH.md`](RALLY_QUALITY_RESEARCH.md) §5.6. **An owner-commissioned 31-reference review of the quality report** *validated* the methodology and surfaced these **future** levers (research inputs, not active work, not new claims; verify citations before any submission):

71 1. **"First-mile" paper framing (highest-value):** stroke-classifiers (BST/DSTA) + the big datasets (ShuttleSet ~36k strokes, Fine-Badminton ~27k actions) all **assume** a pre-trimmed rally clip already exists ? we solve the ignored prerequisite. Pitch the paper as **"the missing untrimmed-video preprocessing pipeline that lets BST/MonoTrack run on raw footage."**

72 2. **Ground R2 in Action-Spotting / Precise-Event-Spotting** (SoccerNet/TenniSet/T-DEED) ? turns R2 into a defensible publishable contribution.

73 3. **Add domain baselines for any submission:** MonoTrack + BST on ShuttleSet / Fine-Badminton (we have only heuristic + Gemini today).

74 4. **Prefer TemporalMaxer over ActionFormer** for the M0.4 scorer swap (parameter-free max-pool, ~2.8x fewer GMACs, small-N robust) ? see "M0.4 next increment" below.

75 5. ****Court polygons ? pose ? the rally-END lever:**** masking on-court players can
 resolve **occluded end-of-rally states** (the #1 remaining gap vs Gemini), not just spectator
 noise.

76 6. ****Name the eval?serve lesson "pipeline distribution skew"**** ? a citable paper
 "lesson."

77

78 ****Discipline reminder:**** owner approval before starting any owner-gated item; one PR per
 step; babysit the PR (poll/fix/reply) until merge observed, **then** advance.

79

80 **## Prioritized next steps**

81 1. ****[owner, in progress] GROW THE GOLDEN CORPUS**** ? the single highest-leverage move;
 every lever below feeds on it.

82 Priority order for new videos: ****(a) multi-court badminton**** (the target distribution
 AND where the ship target

83 lives), ****(b) ?3 matches per new sport ? table tennis first**** (WASB transfers at F1
 0.909; pickleball labels are

84 corpus-gold but not trainable until a clean MIT/Apache ball detector exists), (c)
 capture variety per

85 ``VIDEO_RECORDING_GUIDELINES.md``; ****adult sessions only for the training pool****
 (consent guardrail). Per video:

86 label ? ``curate import-local --normalize --license Proprietary --fps <true fps>`` ?
 re-run the Gen-0 harness.

87 The paired sign test needs n: at n=10 with 9 wins, $p \geq 0.011$? significance is
 reachable this month.

88 2. ****[owner decision] Set the ship-gate target**** ? now informed: floor = heuristic 0.480
 (necessary), ****multi-court**

89 reference = wrapper 0.66** (the number that makes the owned model worth serving),
 clean-footage ceiling = 0.93

90 (not the target ? Gemini serves that tier). Suggested shape: "LOVO F1@tIoU0.5 ? 0.70
 on multi-court folds with

91 paired-sign $p < 0.05$ vs heuristic" ? owner to ratify/adjust.

92 3. ****M0.4 next increment (\$0/local):**** swap the Gen-0 LogReg for a lightweight temporal
 head (MIT, minutes on the

93 2070S ? NOT the heavyweight Gen-2 stack) inside the existing harness; per-video
 threshold calibration (the

94 LOVO-vs-oracle gap is visible per fold in the harness output). ****Prefer**
TemporalMaxer** (parameter-free

95 max-pool, ~2.8x fewer GMACs, small-N robust) over ActionFormer/ASFormer per the
 2026-06-17 external review

96 (``RALLY_QUALITY_RESEARCH.md`` §5.6) ? better fit for the cheap/local/fast mandate.

97 4. ****Multisport thread:**** ingest Roora pickleball/TT CSVs (``import-local --normalize``);
 merged-prototype LOVO across

98 badminton+TT once TT has ?3 matches; identify a clean pickleball ball detector. ? TT-
 on-WASB serving stays gated

99 by C3 (provenance multiplies per sport).

100 5. ****PMF validation (cheaper than any engineering month):**** C13 academy interviews (+
 the two added questions:

101 "multi-court footage nobody watches?" / "pay per-match to index it?"); camera pilot
 at the warmest 2-3 academies

102 (consent at capture; demo served via the Gemini path or human-polish ? the reel
 pattern exists:

103 ``output/MahadevpuraSingles_highlights.mp4``). ****Field kit COMPLETE (2026-06-12, owner**
 ratified the

104 physical-visit approach; reframed product-feedback-first per owner spec):**

105 ``docs/PRODUCT_FIELD_ASSOCIATE_FLYER.md`` (the recruiting flyer ? owner fills ?/contact
 placeholders and recruits) +

106 ``docs/FIELD_VISIT_PLAYBOOK.md`` (the hired associate's visit playbook: coach + student
 script, demo-after-Q6

107 rule, ground rules, Phase-2 record-and-show-back stretch goal) +
``docs/FIELD_VISIT_ANSWER_SHEET.md`` (per-visit

108 capture form incl. demo-reaction + student blocks) + ``docs/PMF_FIELD_SIGNALS.md``
 (G/A/R anchors + ****pre-registered**

109 decision rules** ? PROCEED / PIVOT A-C / FALSIFIED thresholds fixed before the first
 visit).

110 6. ****Housekeeping:**** ~~rotate the Gemini API key~~ ****DONE 2026-06-10**** (owner-verified

working; confirm old-key

```

111     deletion in AI Studio); **rename repo off "badminton" (owner action,
112     in BACKLOG); collector=SoR confirmation; off-peak niceties: none pending ? the
battery completed.
113
114     ## Pointers / artifacts
115     - Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Topology: ADR-008
(`docs/archives/decisions/DECISIONS.md`) ·
116     Study: `docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md` · Quality history:
117     `docs/archives/quality-iterations/2026-06-multivideo-lovo/`.
118     - Harness: `training/gen0/harness.py` (one command ? train ? LOVO ? gate ? artifact
bundle; weights gitignored,
119     manifest reviewable). Featurizer: `backend/features/rally_seq.py` (FEATURE_SCHEMA v1).
120     - Corpus: collector DB (system-of-record, 6 matches/262 rallies) · eval manifest
`output/human_lovo_manifest.json`
121     (gitignored) · durable trajectories + Roora files: `C:\Users\avidu\Projects\Annotation
Setup\Collect\Trajectories\` ·
122     labels: `...\Collect\Golden Labelled\`.
123     - Gemini ops learnings: serial uploads (`indexing.ai_max_workers=1`) + chunk re-encode
124     (`indexing.ai_chunk_scale_width=1280`) + 6-attempt exponential generate retry ? a
cold/prepaid project burst-throttles
125     with a MISLEADING "credits depleted" 429; ramp gradually.
126     - ? No owned-model implementation beyond greenlit M0.x until owner greenlights the next
milestone. Committed next
127     PLATFORM slice = async job model (single-tenant).
128

```

5. user (2026-06-19T17:18:18.485Z)

```

1     - [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Latest (v8) vendor
labeling guidelines mirrored from the collector's `docs/annotation/`, with the fusion-driven
court-calibration / singles-doubles additions marked `? ADDED (fusion)`.
2     - [RALLY_QUALITY_RESEARCH.md](RALLY_QUALITY_RESEARCH.md) ? *** Foundation doc** for
closing the local-vs-Gemini rally-detection gap: the measured gap + the code-grounded over-merge
diagnosis ("shuttle moving ? rally"), an inventory of the eval/experiment harness + its gaps,
the strengthened evaluation framework (segmental Edit/F1@k that make over-segmentation first-
class, golden-corpus + active learning, Gemini-agreement as a *shadow* signal), an
**adversarially-verified, cited external-research synthesis** (TAL/TAS boundary heads, racket-
sports rally-state cues, small-data, eval rigor), a **prioritized cheap?durable roadmap with
falsifiable success criteria**, and (§7) a **tracked project plan** ? work-item checklist,
sequencing, the single quality **number** attributed to the effort, and a **definition of done**
(project completes when all tiers land + the number is recorded). The plan that the
[QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) log executes against; ties together
EXPERIMENT_HARNESS / MULTI_SIGNAL_FUSION_PLAN / ANALYZERS_AND_RECONCILERS /
OWNED_MODEL_IMPLEMENTATION_PLAN.
3     - [RALLY_DETECTION_GAP_CLOSURE.md](RALLY_DETECTION_GAP_CLOSURE.md) ? **Tracked execution
doc (2026-06-18)** for closing the remaining rally-detection accuracy gap **per video, before
launch**, driven by the growing golden corpus. First ground-truth clip (GX010141) reframed it:
boundaries are tight (R2/R3/R4 holds) ? the gap is **detection**: (G1) recall (local & Gemini
both miss ~half on hard clips, and *different* rallies ? complementary), (G2) precision (local
over-fires), (G3) **rally-END over-extension on the quick-return-after-point** (shuttle still
moving ? needs a *player-state* signal). Levers (owner-gated): player-kinematics + shuttle-in-
hand / sustained-invisibility end-rule with reverse-timeline backtrack, `rally_gate` Gate A
wiring, complementary fusion. **Peer** of `RALLY_QUALITY_RESEARCH.md` (execution tracker, not
research); §7 is the source of truth.
4     - [PROMINENT_COURT_DETECTION.md](PROMINENT_COURT_DETECTION.md) ? **Design + experiment
(2026-06-19, default-OFF, owner-gated flip)** for the **multicourt G2 over-fire**: identify the
**prominent (foreground) court** (the one covering the majority of the frame) and keep only
rallies whose shuttle lives inside it ? re-defining the shuttle signal `(x,y,visible)` ?
`(x,y,visible-AND-inside-prominent-court)` *before* windowing, so a background/adjacent-court
shuttle never forms a rally. 2D floor polygon (MVP) ? **pseudo-3D play-volume** (refinement, so
high clears aren't clipped). Born from the GX010142 rally-#1 background-court FP; reuses
`fusion_features.point_in_polygon` + `FusionConfig.court_polygon`. §8 progress tracker is the
source of truth.
5     - [RALLY_DETECTION_QUALITY_REPORT.md](RALLY_DETECTION_QUALITY_REPORT.md) ? **Technical-

```

paper checkpoint (2026-06-17)** of rally-detection quality: the honest measured state (over-segmentation milestone tolerance-F1 0.09170.323 at n=11, p=0.002; merge-rate 0.69470.150; first real-footage ground-truth at-par with START at Gemini parity), the task-faithful R2 metric, a detailed Gemini comparison, next steps + expected wins, areas yet to be explored, and a publishability / path-to-submission verdict. Renders to a polished PDF; the externally-shareable "toward a paper" summary of `RALLY\_QUALITY\_RESEARCH.md` + `QUALITY\_ITERATIONS.md`.

6 - EVALUATION.md ? How we measure and improve quality.

7 - EXPERIMENT_HARNESS.md ? Design (P1 SHIPPED) for A/B + shadow + SxS experimentation on rally-detection variants with **zero production regression** ? experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback = flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`, `metrics`); implemented in `backend/eval/experiment.py`.

8 - QUALITY_ITERATIONS.md ? **Living log** of rally-detection quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-chronological record; terminal-state snapshots get archived under `archives/quality-iterations/`.

9 - ANALYZERS_AND_RECONCILERS.md ? Design (owner-gated, NOT started): a **signal-fusion framework** for rally detection ? producers emit confidence-scored **signals** at three temporal **scopes** (frame detectors / window analyzers / session analyzers), a learned **scorer** weights them, and an auditable **report-first reconciliation** pass lints the decision against the evidence. The spine: **no special-casing in the decision path** (signals?scorer, never `if/else`). Worked examples: a service-fault analyzer, a warm-up/practice span detector, and consuming a per-frame shuttle-state detector. A peer to `EXPERIMENT\_HARNESS.md` and `MULTI\_SIGNAL\_FUSION\_PLAN.md` (ADR-007 modular rally layer).

10 - PERSONALIZATION_PLAN.md ? Design (owner-adopted **hybrid**; Phase-0 landing) for the **per-user personalization** feature on a generalization-first baseline: a thin per-user **tuning** layer gated by the honest small-N statistics (a **min-N promotion floor** + a **structural-guard exemption**, never a service gate) + a **contribution flywheel** (free-tier videos pool into the owner-trained baseline; **"a tool that gets better the more you help it"**). Records the **locked owner decisions** (identity, storage, `min\_n` semantics, Gate-A default, free-tier pooling). Built so far: per-user promotion gate (#141) + held-out-USER learning curve (#142).

11
12 - GOLDEN_REGRESSION_FIXTURES.md ? **Design** (owner-gated, NOT started): **video-free, non-attributable regression fixtures** so a clean clone runs the rally-seg suite with **no video/GPU/DB**. Two layers ? the post-detector freeze + **dual gates** (characterization + quality-floor) reusing the `backend/eval/` stack, and the **SEALED-FIXTURES** privacy filter (de-attribution-at-source, tiered public/key-gated, shuttle-only public tier). Carries a cited **IN/US/EU legal memo** (derived-data IP/privacy), a threat model, and a **\$7 progress tracker + definition-of-done**. Extends GOLDEN_DATA_SHARING.md.

13
14 **## Product / platform plans**
15 - DESKTOP_APP_PLAN.md ? **Tracked scoping doc** (2026-06-18, owner-gated, `[design]` ? NOTHING built): **package the local pipeline as a cross-platform desktop app** for non-technical players ? plug in the camera's USB/SD card, click once, get back **only the highlight reels**, fully on-device (no upload, no original-video bloat). The centerpiece is the **de-WSL native-compute port** (run WASB natively per-OS: Linux/Windows CUDA · macOS MPS/CPU ? drop WSL2), and its **make-or-break unknown**: is CPU acceptable on a typical **no-GPU** enthusiast laptop (WASB is ~3.4x real-time on a laptop 2070S)? \$7 tracker = **P0 compute-feasibility spike** ? P1 native detector ? P2 app shell ? P3 packaging/installers. Realizes the **L0 fully-local** tier of VIDEO_LOCALITY_MODEL.md; is the **LocalDesktop` ComputeBackend** of PLATFORM_ARCHITECTURE.md §3.2; the smaller/export-clean owned model (OWNED_MODEL_IMPLEMENTATION_PLAN.md) is one answer to the compute crux.

16
17 **## Practical / ongoing**

18 - TRACKNET_WSL_SETUP.md

19 - Golden set docs (`GOLDEN\_SET`)

20 - Evaluation harness and calibration drivers (in `backend/eval/`)

21 - GOLDEN_DATA_SHARING.md ? How the GPU box and CPU dev/agent sessions share golden artifacts (trajectory CSVs, `\_golden\_features.json`) via a OneDrive folder; the GPU-extract ? CPU-litmus split. See **##175**.

22

```

23     ## Conventions & templates
24     - [PROJECT_DOC_TEMPLATE.md](PROJECT_DOC_TEMPLATE.md) ? Reusable skeleton for a **tracked
project doc**: status header · decisions-locked · progress tracker (??/?) · definition-of-done
· archived on completion. Standardizes the emergent pattern in `RALLY_QUALITY_RESEARCH.md` §7
and `GOLDEN_REGRESSION_FIXTURES.md`. Use for any substantial design/project work.
25
26     ## Historical documents
27     Completed plans, old research, ADRs, and checkpoints live in [`archives/`](archives/).
28
29     - [CODE_AUDIT_AND_TEST_HARDENING.md](CODE_AUDIT_AND_TEST_HARDENING.md) ? **COMPLETED**
(2026-06-14; T1?T5 + full loop + verification per #162). Living record of the June 2026 plumbing
/ contract / architectural audit, 626-test infrastructure, positive test plans for the 5
identified issues, multi-layer regression protocol (including ablation/experiment layer), and
pre/post-change verification checklist. See `docs/archives/CODE_AUDIT_AND_TEST_HARDENING-
COMPLETED-2026-06-14.md` and `docs/archives/checkpoints/2026-06-14-hardening-loop-complete.md`
for the final archived state. (Top-level is now pointer stub per review; full content in
archives/.)
30
31     See `archives/README.md` for the umbrella policy and `archives/past_projects/README.md`
for classification of small tech debt / starter projects.

```

6. assistant (2026-06-19T17:18:20.588Z)

Now let me read all six docs in my batch.

7. user (2026-06-19T17:18:22.076Z)

```

1     # Field Visit Answer Sheet ? one per academy (fill the same day)
2
3     > Print one per planned visit. Write THEIR words, not yours. A clear "no, because?"
4     > is a *good* sheet. Companion script: `FIELD_VISIT_PLAYBOOK.md`.
5
6     ## Visit basics
7
8     | | |
9     |---|---|
10    | Date / time | |
11    | Academy name | |
12    | Area (neighbourhood) | |
13    | # courts in the hall | |
14    | Spoke to (name + role) | |
15    | Contact (phone/WhatsApp) ? *only if offered* | |
16    | How long the conversation ran | |
17
18    ## The seven questions ? their answers (verbatim where possible)
19
20    **Q1 ? Do they record sessions/matches today? How (phone / tripod / CCTV / hired)?**
21
22    >
23
24    **Q2 ? ? Footage sitting around that nobody watches? Roughly how much, and from what?**
25
26    >
27
28    **Q3 ? Who would watch rally clips/highlights (coach / players / parents)? When?**
29
30    >
31
32    **Q4 ? What do they use today (Machaxi / VISIST / Game Theory / Stupa / BadPro+ / none)?
33    What's good, what's missing, what does it cost them?**
34
35    >
36
37    **Q5 ? ? Would they pay per match for rally cuts + reel + summary? What number did
38    THEY say (?)? What was their face doing?**

```


39

40 Their number came (tick one): ? BEFORE the ladder (they volunteered it) .

8. user (2026-06-19T17:18:22.491Z)

```
1      # Field Visit Playbook ? academy conversations (KhelSutra)
2
3      > **Recruiting?** Hand candidates
[`PRODUCT_FIELD_ASSOCIATE_FLYER.md`](PRODUCT_FIELD_ASSOCIATE_FLYER.md)
4      > ? that's the role description. **This document is the working playbook the hired
5      > associate gets on day one:** the conversation script, the ground rules, and the
6      > Phase-2 (recording / show-back) procedure.
7
8      **Duration:** ~3?4 weeks, part-time (flexible hours, mostly afternoons/evenings when
academies are active)
9      **Location:** Bengaluru ? in-person visits to badminton academies
10     **Compensation:** ? ___ per completed academy visit + travel reimbursement; Phase-2
sessions paid extra ? ? ___ per session *(rates filled at hire; same numbers as the recruiting
flyer)*
11     **Reports to:** Avi (phone/WhatsApp: ___ )
12
13     ---
14
15     ## The gig in one paragraph
16
17     KhelSutra takes ordinary practice/match video from a badminton academy and finds the
rallies automatically, builds
18     highlight reels, and generates a simple summary (rallies found, actual play time); we're
building a private web app
19     that a few invited academies will be able to try. We need honest, on-the-ground feedback
from real academies ? **how
20     they work with video, whether this actually helps them, and what would make it genuinely
useful**. **Your job: visit
21     10?15 good academies, have a friendly 20?30 minute conversation with the owner/head
coach (and, with permission, a
22     few adult (18+) students) using the script below, and write down what they actually
say.** No selling. No tech knowledge needed. For interested academies there's a Phase-2 follow-
up
23     (a recording session ? then you return and show them THEIR OWN footage as highlights)
with additional compensation.
24
25     ---
26
27     ## Phase 1 ? The conversation (what you actually do)
28
29     **Who to talk to:** the academy owner or head coach ? the person who decides what the
academy spends money on.
30     If you get a junior coach, be friendly, gather what you can, and ask when the owner is
available.
31
32     **Opening (keep it this honest):**
33     > "Hi, I'm helping a Bengaluru engineer who's building an AI tool for badminton
academies ? it takes normal practice
34     > video and automatically cuts it into rallies and highlights. He asked me to talk to a
few serious academies to
35     > understand how you work with video today. Could I ask you a few questions? About 15?20
minutes, and we're not
36     > selling anything today."
37
38     **The questions (ask in this order; the starred ones matter most):**
39
40     1. Do you ever video-record sessions or matches here? How ? phone, tripod, CCTV, a hired
person?
```

9. user (2026-06-19T17:18:23.441Z)

```

1      # Field & Product Outreach Associate ? KhelSutra (Sports-Tech, Bengaluru)
2
3      **Help a Bengaluru sports-AI product meet its first real users.**
4
5      | | |
6      |---|---|
7      | **Duration** | ~3?4 weeks, part-time (flexible; mostly afternoons/evenings when
academies are active) |
8      | **Location** | Bengaluru ? in-person visits to badminton academies |
9      | **Pay** | ? ___ per completed academy visit + travel reimbursement; recording sessions
(stretch goal below) paid extra ? ? ___ per session |
10     | **Reports to** | Avi (phone/WhatsApp: ___ ) |
11
12     ---
13
14     ## The product, in plain words
15
16     **KhelSutra** turns ordinary practice/match video from a badminton academy into
17     ready-to-use pieces: it finds the rallies in the recording automatically, cuts
18     **highlight reels**, and produces a **simple summary** (how many rallies, how much
19     actual play time). It works on video from a phone on a tripod ? no special
20     cameras, no hardware to install. We're building a private web app that a few
21     invited academies will be able to try.
22
23     ## The job, in one paragraph
24
25     Visit good badminton academies across Bengaluru and have honest conversations with
26     **coaches AND students**: how do they work with video today, would something like
27     this actually help them, and **what would make it genuinely useful for them** ?
28     including what's missing or wrong with the idea. You work from a provided
29     conversation guide, write down what people *actually say* (the negative answers
30     matter most), and debrief with Avi after each visit. **You are gathering feedback,
31     not selling.** No tech background needed.
32
33     ## The stretch goal (extra pay)
34
35     For academies that are interested, you coordinate a simple recording session ?
36     a phone on a tripod, one-page checklist provided. We run that video through
37     KhelSutra, and you go back and show them **their own court, their own players, as
38     highlight reels and a report**: "this is what it looks like for you." Then you
39     write down exactly what they say. There is no stronger feedback than that moment.
40

```

10. user (2026-06-19T17:18:24.376Z)

```

1      # C13 Field Signals ? what academy visits must tell us (and when to change course)
2
3      **Status:** ACTIVE ? the decision layer for the C13 beachhead validation
4      (owner-ratified direction, 2026-06-12: *PMF validation is handled by physically
5      visiting academies and reading signals on the ground*).
6      **Companions:** [`PRODUCT_FIELD_ASSOCIATE_FLYER.md`](PRODUCT_FIELD_ASSOCIATE_FLYER.md)
7      (the recruiting flyer candidates see) .
8      [`FIELD_VISIT_PLAYBOOK.md`](FIELD_VISIT_PLAYBOOK.md)
9      (the hired associate's visit playbook: script + ground rules + Phase-2) .
10     [`FIELD_VISIT_ANSWER_SHEET.md`](FIELD_VISIT_ANSWER_SHEET.md) (the per-visit capture
form) .
11     [`PMF_MARKET_RESEARCH.md`](PMF_MARKET_RESEARCH.md) (the hypotheses under test, esp. §6).
12
13     **Why this doc exists:** the playbook scripts the questions and the answer sheet
14     collects what was said; this doc fixes ? **before the first visit** ? what those
15     answers *mean* and what totals trigger which decision.
16     Pre-registering the thresholds is the honesty mechanism: it stops us from reading
17     15 ambiguous conversations as whatever we hoped to hear.
18
19     ---

```

```

20
21     ## 1. The hypotheses under test (each maps to script questions)
22
23     | ID | Hypothesis (from PMF research) | Falsified looks like | Script Q |
24     |---|---|---|---|
25     | H1 | **Recording behavior exists** ? academies already produce fixed-angle / CCTV /
multi-court footage without us | nobody records; "we'd have to start for you" | Q1 |
26     | H2 | **Dead-footage pain** ? recorded footage piles up unwatched (the indexing wedge)
| footage is small, or already gets watched | Q2 ? |
27     | H3 | **An audience exists** ? coach, player, or parent would actually consume rally
clips/reels | shrug; "who has time to watch clips" | Q3 |
28     | H4 | **Whitespace** ? incumbents (Machaxi, VISIST, Game Theory, Stupa, BadPro+ ? the
India set is field-sourced from the 2026-06-10 competitive audit; PMF $4's web table predates
it) leave a gap we fit | "we use X, it's fine, fully solved" | Q4 |
29     | H5 | **Willingness to pay per match** at a price that clears unit economics | "should
be free" / <?100 everywhere | Q5 ? |
30     | H6 | **Multi-court chaos is a felt problem** (our moat: indexing messy BYO multi-court
footage) | multi-court video already useful to them | Q6 |
31     | H7 | **Behavioral commitment** ? they'll let us record on their court (Phase-2 door) |
polite no across the board | Q7 |
32
33     Unit-economics floor behind H5: serving COGS on the Gemini path is ~$0.30?0.60 per
34     hour of video (? ?25?50/hr; 2026-06-10 quality table). A 2?3 hr session costs
35     roughly ?75?150 to process, so **?300/match has real margin, ?100/match is
36     break-even-ish at retail** ? that's why the script's anchor ladder is ?100/300/500.
37
38     ## 2. Per-visit scoring ? GREEN / AMBER / RED anchors
39
40     **Scoring is done by the owner (or a Claude session) at transcription ? the

```

11. user (2026-06-19T17:18:25.320Z)

```

1     # PMF & Market Study ? AI Highlight-Reel Tool for Self-Recorded Racket Sports
2
3     > **Research snapshot ? 2026-06-07.** Figures are cited as-of with source + year;
4     > validate live before relying on them for investment/pricing decisions. This is the
5     > market-research-of-record behind the PMF section of the living
6     > [ `COMMERCIALIZATION.md` ](COMMERCIALIZATION.md) tracker ($6). Product framing:
7     > a closed-source hosted API/SaaS (plus local mode) that ingests long-form,
8     > self-recorded, **static-camera, full-court amateur** racket-sports video and
9     > auto-detects rallies and compiles highlight reels. Primary sport: badminton;
10    > sport-generic architecture roadmapped to tennis, pickleball, table tennis, cricket.
11
12    > **Confidence convention.** Figures from named industry reports, federations, or
13    > company pricing pages are tagged with source + year. Inferred items are tagged
14    > **(estimate)** / **(low confidence)**. Emerging-category market sizes vary widely
15    > between vendors; ranges are shown rather than false precision.
16
17    ---
18
19    ## 1. Executive summary
20
21    The product sits at the intersection of three favorable trends: (a) a fast-growing
22    **sports analytics market** (~USD 5?6B in 2025, ~18?22% CAGR to USD 24?36B by ~2034
23    depending on vendor); (b) a structural shift to **amateur self-recording** on action
24    cameras and phones (action-camera market ~USD 6.5B in 2024, sports the largest
25    segment); and (c) **strong racket-sport participation** ? 220M+ badminton players,
26    106M tennis players (2024), and a near-vertical pickleball curve (~19.8M US players in
27    2024, +45.8% YoY).
28
29    The central tension is the **static-camera constraint**: the tool works best on fixed,
30    full-court amateur footage and *fails* on broadcast/panning footage. This is both a
31    moat (most competitors target team sports / broadcast) and a filter on the addressable
32    user ? the customer must already set a camera on a tripod behind the court. That
33    behavior is concentrated in **coaches/academies, clubs, and serious enthusiasts**, not

```

34 the median casual player.
35
36 ****Recommended beachhead (strongest single):**** ****badminton coaching academies and**
37 **serious club/competitive players in India**** ? a large, rapidly formalizing academy
38 ecosystem (Khelo India tailwind), genuine fixed-angle recording behavior, existing
39 willingness to pay for coaching, an underserved tooling gap (incumbents are
40 team-sports/broadcast), and a direct-but-immature competitor field. The product is

12. user (2026-06-19T17:18:25.877Z)

```
1      # KhelSutra ? Product Overview & Early-User Guide
2
3      *What it does today, how to get the best out of it, and where it's going.*
4      *For field associates and for invited alpha / beta academies, coaches, and players.*
5
6      ---
7
8      ## In one line
9
10     **KhelSutra turns ordinary badminton video into useful pieces ? automatically.** Put a
phone or GoPro
11     on a tripod behind the court, record a session, and KhelSutra automatically finds the
rallies, stitches
12     a **highlight reel**, and gives you a **simple summary** (how many rallies, how much
real play time).
13     No special hardware, no editing, no manual scrubbing.
14
15     ---
16
17     ## The problem we're solving
18
19     Academies and serious players already record a lot of video ? and then almost never
watch it. Hours of
20     footage pile up on phones and memory cards because cutting it into useful clips by hand
is tedious and
21     nobody has the time. The good rallies, the moments worth reviewing, stay buried.
22
23     Most academy halls also run **several courts at once**, so a single recording is a busy,
noisy scene
24     that ordinary video tools choke on. **Pulling the real rallies out of busy, multi-court
footage is the
25     hard problem KhelSutra is built to crack** ? it already works very well on clean,
single-court
26     recordings, and the multi-court case is exactly what we're sharpening now.
27
28     ---
29
30     ## Who it's for
31
32     - **Coaching academies & coaches** ? turn raw session footage into review-ready clips
and reels for
33     teaching, without an editor.
34     - **Serious players** ? get clips of your own rallies to review, keep, and share.
35     - Today's focus is **badminton academies in Bengaluru** ? that's who we're building with
first.
36
37     ---
38
39     ## What it does **today**
40
```

13. assistant (2026-06-19T17:18:31.489Z)

All six docs read. The four field-kit docs and PMF_FIELD_SIGNALS are clearly ACTIVE/in-progress (PMF validation track). KHELSUTRA is a fresh untracked product overview. PMF_MARKET_RESEARCH is

a research snapshot still cited as live. Now let me cross-check citations and inbound references for each.

14. user (2026-06-19T17:18:32.554Z)

Found 3 files
docs\NEXT_STEPS.md
docs\PMF_FIELD_SIGNALS.md
docs\FIELD_VISIT_PLAYBOOK.md

15. user (2026-06-19T17:18:33.038Z)

Found 3 files
docs\NEXT_STEPS.md
docs\PMF_FIELD_SIGNALS.md
docs\FIELD_VISIT_ANSWER_SHEET.md

16. user (2026-06-19T17:18:33.982Z)

Found 3 files
docs\NEXT_STEPS.md
docs\PMF_FIELD_SIGNALS.md
docs\FIELD_VISIT_PLAYBOOK.md

17. user (2026-06-19T17:18:34.437Z)

Found 1 file
docs\NEXT_STEPS.md

18. user (2026-06-19T17:18:35.402Z)

Found 7 files
docs\MULTI_SIGNAL_FUSION_PLAN.md
CLAUDE.md
docs\PLATFORM_ARCHITECTURE.md
docs\COMMERCIALIZATION.md
docs\PMF_FIELD_SIGNALS.md
docs\I18N_PLAN.md
docs\archives\checkpoints\2026-06-strategy-foundation\README.md

19. user (2026-06-19T17:18:35.638Z)

No files found

20. assistant (2026-06-19T17:18:40.578Z)

KHELSUTRA_PRODUCT_OVERVIEW.md has zero inbound references (it's the untracked new file from git status). Let me check the status headers / lifecycle of PMF_MARKET_RESEARCH and KHELSUTRA more closely, plus confirm the field docs have no archived copies and check PMF_FIELD_SIGNALS status.

21. user (2026-06-19T17:18:44.095Z)

=== archives glob for batch stems ===
=== none above means no archived copies ===

=== KHELSUTRA git tracked? ===
(empty = untracked)

=== how CLAUDE.md cites PMF_MARKET_RESEARCH ===
13:- **Strategy (post-scaffolding):** `docs/COMMERCIALIZATION.md`,
`docs/PMF\_MARKET\_RESEARCH.md`,

22. assistant (2026-06-19T17:18:48.144Z)

Confirmed. Now let me check the tail of PMF_FIELD_SIGNALS for decision-rule status, and scan PMF_MARKET_RESEARCH for any "superseded/completed" markers, plus check the rest of KHELSUTRA for any status header.

23. user (2026-06-19T17:18:50.110Z)

52:variance is the core failure mode; (4) badminton's strongest markets (India, Indonesia, 87:| Badminton ? Indonesia | National sport; ~1M club players | industry (est.) | 174:### #3 ? Badminton academies/clubs, **East/SE Asia (China, Indonesia, Malaysia, Japan, Korea) + Denmark** 198:4. **Monetization mismatch in best-participation markets** (India/Indonesia/China are

24. user (2026-06-19T17:18:51.429Z)

1:# KhelSutra ? Product Overview & Early-User Guide 4:*For field associates and for invited alpha / beta academies, coaches, and players.* 8:## In one line 17:## The problem we're solving 30:## Who it's for 39:## What it does **today** 57:## How to use it (alpha) 70:## How to get the best out of it (capture matters) 86:## The part that makes KhelSutra different: **it gets sharper from real academy footage** 111:## Where it's going ? toward an **AI assistant coach** 130:## What feedback helps us most (for alpha users) 144:## What we're focused on over the next ~4 weeks 156:## What KhelSutra is **not** (so expectations stay honest) 166:## Quick answers to common questions

25. assistant (2026-06-19T17:19:01.557Z)

I have enough evidence for all six docs. Summary of findings:

- **All four field-kit docs + PMF_FIELD_SIGNALS**: The PMF/field-visit validation track is explicitly LIVE and in-progress. NEXT_STEPS.md §5 lists "Field kit COMPLETE (2026-06-12)" but that means the kit was *built* and is staged for owner to fill placeholders and recruit ? the validation work (the actual visits) has NOT happened. The docs contain blank capture forms, unfilled ?/contact placeholders, and pre-registered decision rules awaiting first visit. These are active operational instruments, not terminal-state work. No archived copies exist. All KEEP_LIVE.
- **PMF_FIELD_SIGNALS.md**: status header literally says "ACTIVE ? the decision layer for C13 beachhead validation." KEEP_LIVE.
- **PMF_MARKET_RESEARCH.md**: a "Research snapshot 2026-06-07" but cited as a LIVE strategy source in CLAUDE.md ("Strategy (post-scaffolding)") and is the market-research-of-record behind COMMERCIALIZATION.md §6. Strong inbound refs (7 files). Not terminal. KEEP_LIVE.
- **KHELSUTRA_PRODUCT_OVERVIEW.md**: brand-new untracked file (in git status as `??`), describes the product "today" for alpha/beta users, has a "where it's going" roadmap ? a living product overview, zero inbound refs, no status header. New/active, not terminal. KEEP_LIVE.

None are in a terminal DONE/REJECTED/SUPERSEDED state; none have archived copies; none are duplicates.

26. user (2026-06-19T17:19:25.967Z)

Structured output provided successfully

27. assistant (2026-06-19T17:19:34.046Z)

Audit complete. All six docs in my batch are KEEP_LIVE ? none are in a terminal DONE/REJECTED/SUPERSEDED state, none have archived copies, and none are duplicates. Key reasoning:

- The four field-kit docs (`FIELD\_VISIT\_ANSWER\_SHEET.md`, `FIELD\_VISIT\_PLAYBOOK.md`, `PRODUCT\_FIELD\_ASSOCIATE\_FLYER.md`) plus `PMF\_FIELD\_SIGNALS.md` are active operational

instruments for the in-progress C13 PMF validation track. `NEXT\_STEPS.md` §5 says the "field kit COMPLETE (2026-06-12)" but that means the kit was *built* and staged for the owner to fill placeholders, recruit, and run visits ? the validation itself is pending. They still carry blank capture fields and unfilled ?/contact placeholders. `PMF\_FIELD\_SIGNALS.md` literally carries a `Status: ACTIVE` header.

- `PMF\_MARKET\_RESEARCH.md` is a dated 2026-06-07 snapshot but is cited as a LIVE strategy source in `CLAUDE.md` and is the "market-research-of-record" behind the living `COMMERCIALIZATION.md` §6 and the active PMF track (7 inbound referrers).

- `KHELSUTRA\_PRODUCT\_OVERVIEW.md` is a brand-new untracked file (`??` in git status, empty `git ls-files`) ? a living product overview with a forward-looking roadmap, zero inbound refs, no terminal marker.

The full per-doc structured verdicts are in the StructuredOutput call.