

# claude-windows / fa4dc4eb-4705-47fd-b044-4f017b586b15

## Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\fa4dc4eb-4705-47fd-b044-4f017b586b15\subagents\workflows\wf_ac0cbb29-6bf\agent-aa12364503d8455e2.jsonl`  
- SHA-256: `33ededaf48352cdcef3ed83540b65e7a21408fdb4f4257cd5b66de3e5ee4ffc6`  
- Source modified: `2026-06-19T15:53:59+00:00`  
- Imported at: `2026-07-05T16:48:22+00:00`  
- project: `wf_ac0cbb29-6bf`  
- session_id: `fa4dc4eb-4705-47fd-b044-4f017b586b15`
```

## Transcript

### 1. user (2026-06-19T15:52:06.176Z)

You are auditing a badminton rally-detection codebase for DOC/COMMENT STALENESS ? statements that CONTRADICT the current code or data.

HARD RULES:

- Report ONLY factual staleness with CODE EVIDENCE (file:line or symbol) that proves it. Open the cited code and confirm before reporting.
- NEVER invent a fact or a status. If a thing is built-but-default-OFF / unwired / unmeasured, say EXACTLY that ? do not overclaim "in production".
- A "[design]" tag is stale ONLY if the thing is actually BUILT in code (cite the symbol).
- DO NOT propose subjective rewrites, style/tone changes, or improvements. DO NOT relitigate owner-gated DECISIONS (owned-model roadmap, licensing/MIT-Apache stance, product/strategy choices, genuinely-still-design proposals).
- Preserve the document's voice and formatting; make the SMALLEST surgical edit. stale\_text MUST be copied verbatim with enough context to be unique.
- IGNORE docs/archives/\*\* (terminal snapshots) and any memory/ file.

Working dir is the repo root; use Read/Grep/Glob freely.

ADVERSARIALLY VERIFY these proposed staleness fixes for ONE file: docs/NEXT\_STEPS.md.

For EACH proposed fix: (1) open docs/NEXT\_STEPS.md and confirm stale\_text appears VERBATIM (tighten/expand it to an exact, UNIQUE match if needed ? this text will be applied by exact string match, so it must be perfect); (2) re-open the cited code yourself and confirm the claim is genuinely contradicted by the CURRENT code; (3) confirm corrected\_text is accurate, preserves voice, and is minimal; (4) DROP it if it is subjective, an owner-gated decision, already-correct, or you cannot verify the evidence. Output ONLY the survivors (verdict confirm or revise), with stale\_text guaranteed to be an exact unique substring of the file.

PROPOSED FIXES:

```
[  
  {  
    "file": "docs/NEXT_STEPS.md",  
    "stale_text": "***Immediate next step ? the real P3 LOVO measurement, split GPU?CPU** (the dev box is CPU-only: `torch+cpu`/`cv2` present but no GPU, and the corpus/videos/CSVs/labels are not here):",  
    "corrected_text": "***Immediate next step ? the real P3 LOVO measurement** (GPU + torch (`torch+cu124`) + cv2 are present here, but the corpus/videos/CSVs/labels are not ? those remain owner-side):",  
    "why_stale": "Asserts 'the dev box is CPU-only: torch+cpu present but no GPU'. Verified directly: torch 2.6.0+cu124 with torch.cuda.is_available()==True, plus cv2 4.13.0. The environment has a GPU and a CUDA torch build, contradicting 'CPU-only / torch+cpu / no GPU'. (The corpus/videos/labels-not-here part is left intact.)",  
    "severity": "medium"  
  },  
  {  
    "file": "docs/NEXT_STEPS.md",  
    "stale_text": "- ? Dev container no GPU/torch/cv2 ? Gate A, person cue logic, harness, and
```

```

pure fusion paths are testable here; real-video / golden LOVO / ablation confirmation is owner-
side (GPU/WSL machine).",
    "corrected_text": "- ? Gate A, person cue logic, harness, and pure fusion paths are testable
here; the heavier real-video / golden LOVO / ablation confirmation is owner-side (the golden
corpus + labelled videos live on the owner's machine).",
    "why_stale": "Says 'Dev container no GPU/torch/cv2'. Verified directly: torch 2.6.0+cu124
(cuda True) and cv2 4.13.0 both import in this environment, so 'no GPU/torch/cv2' is false. (The
owner-side real-video/golden-LOVO point is preserved, reframed around corpus location rather
than missing GPU.)",
    "severity": "medium"
}
]

```

## 2. assistant (2026-06-19T15:52:38.626Z)

I'll verify these two proposed fixes against the actual file content and the current environment.

## 3. user (2026-06-19T15:52:39.582Z)

```

1      # Next steps (live) ? owned-model + multisport
2
3      **Updated:** 2026-06-13 (**rally-detection quality track**: fusion + experiment-harness
design docs merged,
4      owner-gated ? see the dedicated section below, the desktop-agent pick-up). Prior:
2026-06-10 (audit?M0?Gemini-battery session; later same day the **UI-alpha track opened**:
5      ADR-009 ratified ? KhelSutra Guru / khelsutra.guru, private `khelsutra` repo, local-
first delivery ? see
6      `docs/UI_ALPHA_EXECUTION_PLAN.md`; **PR-1 async jobs (#16) MERGED 2026-06-11, PR #87**;
**PR-2
7      upload/download MERGED 2026-06-11, PR #99** ? live-E2E-verified (3-part chunked upload
MD5 byte-identity +
8      compile?browser-download); next: PR-3 identity. Same day: the **8?9 quality sweep
#90?#98 landed** ? ruff/mypy/
9      coverage-floor CI lanes, SQLite conn-close fix, config unknown-key warning, hybrid-twin
collapse into
10     `trajectory_hybrid.py`, WSL tilde+abspath live-path fixes; suite 388 green; LOVO 0.626
re-verified exact).
11     **Authoritative roadmap:**
12     `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` (+ ADR-008 for topology). **Latest blow-by-
blow:** memory
13     `current-state.md` / `session-handoff.md`. This file = the prioritized "what to do
next," refreshed each session.
14
15     ## ? Next-session / GPU-box pickup (2026-06-15)
16     **M0 in-container quick wins are COMPLETE:** ship-gate provisional blocker (**#170**),
train/eval split guard (**#171**), the **doc-consistency sweep** (**#174** ? corrected 16
drifted docs vs the code), and the **P3 person-feature fusion litmus machinery is already
present + suite-green** (`backend/eval/fusion_compare.py` + `ablation.py`: person-driven ?F1 /
bootstrap CI / paired sign test; the real-golden `skipif` litmus waits only on GPU-extracted
data). ? CI is red repo-wide due to a **GitHub Actions minutes/billing outage** (account-level,
not code) ? work is local-gated + `--admin`-merged until Actions is restored (Settings ? Billing
? Actions minutes / spending limit).
17
18     **Immediate next step ? the real P3 LOVO measurement, split GPU?CPU** (the dev box is
CPU-only: `torch+cpu`/`cv2` present but no GPU, and the corpus/videos/CSVs/labels are not here):
19     1. **[owner / GPU box]** `python -m backend.eval.fusion_golden --manifest
output/human_lovo_manifest.json --out-dir "%USERPROFILE%\OneDrive\rally-annotator\golden-
data\features\2026-06-15-n6" ? produces `_golden_features.json` for the 6 golden videos. (GPU
+ WSL + videos + trajectory CSVs + `rallies.csv` labels are owner-side only.)
20     2. Artifacts land in the **shared cloud folder** `C:\Users\avidu\OneDrive\rally-
annotator\golden-data\` (OneDrive auto-syncs both ways). Convention + workflow:
**[`GOLDEN_DATA_SHARING.md`](GOLDEN_DATA_SHARING.md)**; tracked in **#175**.
21     3. **[CPU dev/agent session]** `python -m backend.eval.fusion_compare --features-dir
<shared>\features\2026-06-15-n6 --record` + `python -m backend.eval.ablation --features-dir

```

```

<shared>\... --granularity group` ? real person-feature ?F1; copy the JSONs into `output/` to
also pass the `skipif` litmus tests (`tests/test_ablation.py::test_litmus_reproduces_gate_a`).
22
23     **Also queued (owner-gated / pending input):**
24     - Ship-gate target calibration ? **#172** (owner decision; needs corpus growth).
25     - Tests reorganization ? blocked on the owner's grouping metric (provide it ? an agent
will do it).
26     - Restore CI ? the Actions outage is account-level (billing/minutes), not a repo/config
bug.
27
28     **Leaving (bulky / M2 ? no start without owner go):** ActionFormer (M0.4), event-index
DB migration (M0.1), Gemini-silver purge (M0.3), cost-to-Gen-2 benchmark + ByteTrack
deprecation, multisport, repo rename, PMF execution.
29
30     ## Where we are (one paragraph)
31     ADR-008 ratified ("repo split driven by productionization, not isolation"): training
lives in-repo (`training/`)
32     behind a CI-enforced import wall; the featurizer is single-sourced
(`backend/features/rally_seq.py`, train==serve);
33     the **Gen-0 harness ships** (`python -m training.gen0.harness --manifest
output/human_lovo_manifest.json --floor
34     eval_baselines/heuristic_lovo_n6.json --version <semver>`) and produced artifact v0.1.0
(LOVO **0.626**, floor passed,
35     sign test 5/6 wins p=0.109 ? honestly not significant, ship=False). The **n=6 owned
corpus is in the collector**
36     (262 rallies, Proprietary, commercial export = owned data only after the ShuttleSet
reclassification). The **Gemini
37     battery is done** (see table) ? the moat is quantified.
38
39     ## The quality table that now drives strategy (IoU 0.5 vs human golden labels)
40     | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |
41     |---|---|---|
42     | Heuristic (velocity windowing) | 0.657 | 0.157 |
43     | **Gen-0 owned head** (LOVO, n=6) | 0.744 | 0.561 |
44     | Two-stage WASB+Gemini refine | 0.83 | ? |
45     | **Gemini wrapper** (`gemini-flash-latest`) | **0.93** | **0.66** |
46
47     **Reading:** clean footage is commoditized (serve it with Gemini for cents). **Multi-
court drops the commodity
48     wrapper to 0.66 ? that is the owned model's ship target** (its FPs are background-game
pickups + dead-time merges,
49     the exact contrast-metric confound). Gen-0 is 0.10 behind with only n=6 ? a corpus-
growth-sized gap, not an
50     architecture-sized one. Two-stage refine = the cost-efficient serving path for LONG
tapes (uploads candidate clips
51     only); wrapper 503-flakiness (observed all session) makes the provider-fallback registry
load-bearing.
52     Baselines tracked in `eval_baselines/` (heuristic_lovo_n6.json,
gemini_wrapper_2026-06-10.json).
53
54     ## Rally-detection quality track (design merged 2026-06-13, owner-gated; implementation
in progress)
55     **This is the track a desktop/local agent is slated to pick up.** Design docs landed (PR
#112 + session). All implementation is **owner-gated** (per explicit in the plan); each step =
ONE PR off `master`.
56
57     **Hardening loop (T1-T5 positive tests + audit + finale) COMPLETE (2026-06-14, PRs
#156/#157/#160/#161/#163 + #165 + final records).** See archived
`docs/archives/CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md` (Living Log with every
pre/post/review) and `docs/archives/HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md` (full rules
+ STATE). Top-level now pointer stubs per #165 review. All 5 items + verification + archive
done; clean slate.
58
59     - **Docs:** `docs/MULTI_SIGNAL_FUSION_PLAN.md` (shuttle + person + audio, one fusion
layer) .

```

```

60     `docs/EXPERIMENT_HARNESS.md` (A/B + shadow + promotion gate, zero-regression) .
61     `docs/ROORA_LABELING_GUIDELINES.md` (v8) . `docs/HOW_RALLY_DETECTION_WORKS.md`
(2-layer mental model).
62     - Key finding (no re-investigation needed): the owner's "held-shuttle counted as a
rally" bug is NOT a missing capability ? `backend/pipeline/detectors/rally_gate.py` already
implements the net-crossing "Gate A" (`apply_gate(..., min_crossings=2)`) that kills
service-feed/held-shuttle windows, but it is only called from eval drivers (and there was no
`min_crossings` knob in `IndexingConfig`). Wiring Gate A into the live segmenter (now via
fusion_hybrid) was the single cheapest, highest-confidence quality win. (P2 FusionSegmenter
specifics: registry `fusion_hybrid`, default-OFF, seeds from substrate, persists
`net_crossings` + optional person features, optional served Gate A/B via config knobs;
adversarially reviewed; suite green. See `tests/test_served_gate_a_regression.py` for the honest
finding that on production windowing Gate A is neutral/negative in some cases, hence kept opt-in
with knob for safety, and the leaderboard revert note.)
63     - Current status (2026-06-14): P1 (person cue extraction to reusable
`PersonDetector`) + P2 (FusionSegmenter + Gate A signal + served handoff gating) + experiment
harness P1 DONE (adversarial review + NO-REGRESSION; suite green). P3 next (this PR + follow-
ups): learned fusion (person features ? harness `compare` LOVO lift on the 6 golden videos;
eager-person-compute optimisation); then P4 audio via hit_detector.
64     - Pick-up order (each = ONE PR off `master`, owner approval first, flag-gated default-
OFF, promote only on measured LOVO):
65         1. (P2 first-step, largely done via fusion_hybrid) Wire Gate A into production served
path + config knob (min_crossings); apply in hybrids at runtime. (Now available opt-in; served
neutral on production windowing per litmus ? kept opt-in with knob for safety.)
66         2. P3 first step (this work item): harness `compare` LOVO litmus for person-on variant
(players_in_court + two_sided + colocation etc. from `fusion_features`) on the 6 golden; add the
eager compute optimisation note/landing. Pure + harness (testable here).
67         3. P3/P4 continuation: audio cue integration + full learned promotion if lift proven.
68     - Discipline (from EXPERIMENT_HARNESS.md): every new cue flag-gated, default
OFF; promote to default only on measured LOVO lift through `run_regression.py` (exit
0/1/3); pinned CONTROL ? rollback = config flip. Harness re-uses ~80% existing code.
69     - ? Dev container no GPU/torch/cv2 ? Gate A, person cue logic, harness, and pure fusion
paths are testable here; real-video / golden LOVO / ablation confirmation is owner-side (GPU/WSL
machine).
70     - External due-diligence review inputs (2026-06-17) ? PICK UP AFTER current ideas are
exhausted; full details + citations in [`RALLY_QUALITY_RESEARCH.md`](RALLY_QUALITY_RESEARCH.md)
$5.6. An owner-commissioned 31-reference review of the quality report validated the
methodology and surfaced these future levers (research inputs, not active work, not new
claims; verify citations before any submission):
71         1. "First-mile" paper framing (highest-value): stroke-classifiers (BST/DSTA) + the
big datasets (ShuttleSet ~36k strokes, Fine-Badminton ~27k actions) all assume a pre-trimmed
rally clip already exists ? we solve the ignored prerequisite. Pitch the paper as "the
missing untrimmed-video preprocessing pipeline that lets BST/MonoTrack run on raw footage."
72         2. Ground R2 in Action-Spotting / Precise-Event-Spotting
(SoccerNet/TenniSet/T-DEED) ? turns R2 into a defensible publishable contribution.
73         3. Add domain baselines for any submission: MonoTrack + BST on ShuttleSet / Fine-
Badminton (we have only heuristic + Gemini today).
74         4. Prefer TemporalMaxer over ActionFormer for the M0.4 scorer swap (parameter-free
max-pool, ~2.8x fewer GMACs, small-N robust) ? see "M0.4 next increment" below.
75         5. Court polygons ? pose ? the rally-END lever: masking on-court players can
resolve occluded end-of-rally states (the #1 remaining gap vs Gemini), not just spectator
noise.
76         6. Name the eval/serve lesson "pipeline distribution skew" ? a citable paper
"lesson."
77
78     Discipline reminder: owner approval before starting any owner-gated item; one PR per
step; babysit the PR (poll/fix/reply) until merge observed, then advance.
79
80     ## Prioritized next steps
81     1. [owner, in progress] GROW THE GOLDEN CORPUS ? the single highest-leverage move;
every lever below feeds on it.
82         Priority order for new videos: (a) multi-court badminton (the target distribution
AND where the ship target
83         lives), (b) ?3 matches per new sport ? table tennis first (WASB transfers at F1
0.909; pickleball labels are

```

```

84 corpus-gold but not trainable until a clean MIT/Apache ball detector exists), (c)
capture variety per
85 `VIDEO_RECORDING_GUIDELINES.md`; **adult sessions only for the training pool**
(consent guardrail). Per video:
86 label ? `curate import-local --normalize --license Proprietary --fps <true fps>` ?
re-run the Gen-0 harness.
87 The paired sign test needs n: at n=10 with 9 wins, p>0.011 ? significance is
reachable this month.
88 2. **[owner decision] Set the ship-gate target** ? now informed: floor = heuristic 0.480
(necessary), **multi-court
89 reference = wrapper 0.66** (the number that makes the owned model worth serving),
clean-footage ceiling = 0.93
90 (not the target ? Gemini serves that tier). Suggested shape: "LOVO F1@tIoU0.5 ? 0.70
on multi-court folds with
91 paired-sign p<0.05 vs heuristic" ? owner to ratify/adjust.
92 3. **M0.4 next increment ($0/local):** swap the Gen-0 LogReg for a lightweight temporal
head (MIT, minutes on the
93 2070S ? NOT the heavyweight Gen-2 stack) inside the existing harness; per-video
threshold calibration (the
94 LOVO-vs-oracle gap is visible per fold in the harness output). **Prefer
TemporalMaxer** (parameter-free
95 max-pool, ~2.8x fewer GMACs, small-N robust) over ActionFormer/ASFormer per the
2026-06-17 external review
96 (`RALLY_QUALITY_RESEARCH.md` $5.6) ? better fit for the cheap/local/fast mandate.
97 4. **Multisport thread:** ingest Roora pickleball/TT CSVs (`import-local --normalize`);
merged-prototype LOVO across
98 badminton+TT once TT has ?3 matches; identify a clean pickleball ball detector. ? TT-
on-WASB serving stays gated
99 by C3 (provenance multiplies per sport).
100 5. **PMF validation (cheaper than any engineering month):** C13 academy interviews (+
the two added questions:
101 "multi-court footage nobody watches?" / "pay per-match to index it?"); camera pilot
at the warmest 2-3 academies
102 (consent at capture; demo served via the Gemini path or human-polish ? the reel
pattern exists:
103 `output/MahadevpuraSingles_highlights.mp4`). **Field kit COMPLETE (2026-06-12, owner
ratified the
104 physical-visit approach; reframed product-feedback-first per owner spec):**
105 `docs/PRODUCT_FIELD_ASSOCIATE_FLYER.md` (the recruiting flyer ? owner fills ?/contact
placeholders and recruits) +
106 `docs/FIELD_VISIT_PLAYBOOK.md` (the hired associate's visit playbook: coach + student
script, demo-after-Q6
107 rule, ground rules, Phase-2 record-and-show-back stretch goal) +
`docs/FIELD_VISIT_ANSWER_SHEET.md` (per-visit
108 capture form incl. demo-reaction + student blocks) + `docs/PMF_FIELD_SIGNALS.md`
(G/A/R anchors + **pre-registered
109 decision rules** ? PROCEED / PIVOT A-C / FALSIFIED thresholds fixed before the first
visit).
110 6. **Housekeeping:** ~~rotate the Gemini API key~~ **DONE 2026-06-10** (owner-verified
working; confirm old-key
111 deletion in AI Studio); **rename repo off "badminton"** (owner action,
112 in BACKLOG); collector=SoR confirmation; off-peak niceties: none pending ? the
battery completed.
113
114 ## Pointers / artifacts
115 - Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Topology: ADR-008
(`docs/archives/decisions/DECISIONS.md`) ·
116 Study: `docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md` · Quality history:
117 `docs/archives/quality-iterations/2026-06-multivideo-lovo/`.
118 - Harness: `training/gen0/harness.py` (one command ? train ? LOVO ? gate ? artifact
bundle; weights gitignored,
119 manifest reviewable). Featurizer: `backend/features/rally_seq.py` (FEATURE_SCHEMA v1).
120 - Corpus: collector DB (system-of-record, 6 matches/262 rallies) · eval manifest
`output/human_lovo_manifest.json`
121 (gitignored) · durable trajectories + Roora files: `C:\Users\avidu\Projects\Annotation

```

```
Setup\Collect\Trajectories\` ` .
122     labels: `...\Collect\Golden Labelled`.
123     - Gemini ops learnings: serial uploads (`indexing.ai_max_workers=1`) + chunk re-encode
124     (`indexing.ai_chunk_scale_width=1280`) + 6-attempt exponential generate retry ? a
cold/prepaid project burst-throttles
125     with a MISLEADING "credits depleted" 429; ramp gradually.
126     - ? No owned-model implementation beyond greenlit M0.x until owner greenlights the next
milestone. Committed next
127     PLATFORM slice = async job model (single-tenant).
128
```

#### 4. user (2026-06-19T15:52:46.326Z)

```
torch 2.6.0+cu124
cuda_available True
cv2 4.13.0
```

#### 5. assistant (2026-06-19T15:52:50.388Z)

Both fixes target the environment-capability claims. I've confirmed the environment has torch 2.6.0+cu124 (CUDA available True) and cv2 4.13.0. Now let me verify the exact text matches and check whether the corpus/data claims (which the corrected text preserves) are accurate.

Let me verify the exact stale\_text strings appear verbatim and are unique.

#### 6. user (2026-06-19T15:52:51.869Z)

18:\*\*Immediate next step ? the real P3 LOVO measurement, split GPU?CPU\*\* (the dev box is CPU-only: `torch+cpu`/`cv2` present but no GPU, and the corpus/videos/CSVs/labels are not here):

#### 7. user (2026-06-19T15:52:52.929Z)

69:- ? Dev container no GPU/torch/cv2 ? Gate A, person cue logic, harness, and pure fusion paths are testable here; real-video / golden LOVO / ablation confirmation is owner-side (GPU/WSL machine).

#### 8. assistant (2026-06-19T15:52:58.860Z)

Both stale\_text strings appear verbatim and uniquely. Now I need to consider a subtlety on fix #1. The corrected\_text changes `torch+cpu` to `torch+cu124`. Let me verify the exact CUDA build string to ensure accuracy.

The torch version reported is `2.6.0+cu124`, so `torch+cu124` is an accurate characterization. Let me also double-check fix #1's corrected\_text doesn't introduce a problem: the original says "split GPU?CPU" in the bold header, and the corrected text drops that. That's consistent with the new claim (no longer a GPU/CPU split since GPU is present here). The data-not-here portion is preserved.

Now let me verify the "corpus/videos/CSVs/labels are not here" claim is still accurate ? that the preserved portion is true. Let me check whether the golden corpus paths exist locally.

#### 9. user (2026-06-19T15:53:01.188Z)

```
total 8
drwxr-xr-x 1 avidu 197609    0 Jun 15 10:57 .
drwxr-xr-x 1 avidu 197609    0 Jun 15 10:57 ..
-rw-r--r-- 1 avidu 197609 1215 Jun 15 09:56 README.md
drwxr-xr-x 1 avidu 197609    0 Jun 15 10:57 baselines
drwxr-xr-x 1 avidu 197609    0 Jun 15 10:57 features
drwxr-xr-x 1 avidu 197609    0 Jun 15 10:57 manifests
drwxr-xr-x 1 avidu 197609    0 Jun 15 10:57 trajectories
---
total 17964
drwxr-xr-x 1 avidu 197609    0 Jun 18 12:26 .
```

```
drwxr-xr-x 1 avidu 197609      0 Jun 15 14:38 ..
-rw-r--r-- 1 avidu 197609 1549493 Jun 17 15:02 Badminton_BXH_2_june15_traj.csv
-rw-r--r-- 1 avidu 197609 838358 Jun 18 11:28 Boxhill_Doubles_june18_traj.csv
-rw-r--r-- 1 avidu 197609 1543847 Jun 16 19:04 GX010128_june12_traj.csv
-rw-r--r-- 1 avidu 197609 797173 Jun 17 18:13 GX010137_june17_traj.csv
-rw-r--r-- 1 avidu 197609 776018 Jun 17 17:23 GX010141_june17_traj.csv
-rw-r--r-- 1 avidu 197609 608149 Jun 18 11:53 GX020094_june18_traj.csv
-rw-r--r-- 1 avidu 197609 701986 Jun 17 14:15 GX030094_june15_traj.csv
---output---
C:/Users/avidu/Projects/badminton-highlight-indexer/output/human_lovo_manifest.json
```

#### 10. assistant (2026-06-19T15:53:04.944Z)

This is an important nuance. The trajectory CSVs and golden-data folder ARE present on this machine (under `C:/Users/avidu/Projects/Annotation Setup/Collect/Trajectories/` and the OneDrive `golden-data/` folder). The `output/human\_lovo\_manifest.json` also exists locally.

This complicates the "the corpus/videos/CSVs/labels are not here" preserved claim. Let me check more carefully what's actually present vs. absent ? specifically the videos and the labelled `.rallies.csv` files.

#### 11. user (2026-06-19T15:53:08.547Z)

```
=== Golden Labelled ===
ls: cannot access 'C:/Users/avidu/Projects/Annotation Setup/Collect/Golden Labelled/': No such
file or directory
=== golden-data/features ===
=== golden-data/trajectories ===
=== videos anywhere? ===
C:/Users/avidu/Projects/Annotation Setup/6 June 2026/Annotation Videos/Sample Long Video -
Baadminton.MP4
C:/Users/avidu/Projects/Annotation Setup/6 June 2026/Annotation Videos/Video 2 - Pickleball.mp4
C:/Users/avidu/Projects/Annotation Setup/6 June 2026/Annotation Videos/Video 3 - Tabletennis.mp4
C:/Users/avidu/Projects/Annotation Setup/Collect/Candidates/Adarsh and Avi (last game).MP4
C:/Users/avidu/Projects/Annotation Setup/Collect/Candidates/Cricket Net.MP4
C:/Users/avidu/Projects/Annotation Setup/Collect/Candidates/GX010055.MP4
C:/Users/avidu/Projects/Annotation Setup/Collect/Candidates/GX010098.MP4
C:/Users/avidu/Projects/Annotation Setup/Collect/Candidates/Mahadevpura Multimatch.MP4
C:/Users/avidu/Projects/Annotation Setup/Collect/Candidates/VeryLargeTestVideo.MP4
=== rallies.csv labels ===
```

#### 12. assistant (2026-06-19T15:53:18.225Z)

The data picture is mixed. Videos and trajectory CSVs ARE present on this machine, but the `Golden Labelled/` directory does not exist, and the `golden-data/features/` and `trajectories/` subfolders are empty. So the labelled `.rallies.csv` files and extracted golden features are genuinely absent here.

This matters for assessing the corrected\_text of both fixes, which assert the data is "not here." Let me look at this precisely:

**\*\*Fix #1 corrected\_text:\*\*** "the corpus/videos/CSVs/labels are not here ? those remain owner-side" ? but videos AND trajectory CSVs ARE here. Only the labels/golden-features are absent. So this corrected\_text would introduce a NEW inaccuracy (claiming videos/CSVs are not here when they are).

**\*\*Fix #2 corrected\_text:\*\*** "the golden corpus + labelled videos live on the owner's machine" ? this is on the owner's machine, but it IS also this machine (the owner's own machine ? the env is the owner's Windows box per the OneDrive/Projects paths). The corrected text reframes around "owner's machine," and the labelled data genuinely isn't extracted/present.

The core verified fact for BOTH fixes is the same and solid: the GPU/torch/cv2 claim is false. That part is unambiguously stale and the primary subject of both fixes. Let me confirm this is the owner's own machine (not a separate "dev container") to validate the framing.

### 13. user (2026-06-19T15:53:20.130Z)

user: avidu  
host: AVIS-MSI  
NVIDIA GeForce RTX 2070 Super with Max-Q Design, 8192 MiB

### 14. assistant (2026-06-19T15:53:31.994Z)

This is confirmed: the machine is `AVIS-MSI` (the owner Avi's MSI laptop, the "2070S" referenced in the doc itself at line 93), with an RTX 2070 Super GPU, torch+cu124, and cv2. The "dev container no GPU/torch/cv2" / "CPU-only" claims are unambiguously false ? this is the owner's GPU box.

Now let me finalize the assessment of each fix:

**\*\*Both fixes:\*\*** The core staleness ? the GPU/torch/cv2 capability claim ? is genuine and well-evidenced. Confirm on that.

**\*\*Fix #1 ? the data-location refinement needs revision.\*\*** The corrected\_text says "the corpus/videos/CSVs/labels are not here." But I verified videos (\*.MP4) and trajectory CSVs (\*\_traj.csv) ARE present on this machine. What's genuinely absent: the `Golden Labelled/` directory (the \*.rallies.csv labels) and extracted golden features. The doc's line 19 specifies the data that's owner-side-only as "GPU + WSL + videos + trajectory CSVs + \*.rallies.csv labels." Since I found videos and traj CSVs present, asserting they're "not here" introduces a new inaccuracy. I'll revise to keep only the verifiably-absent items (labels/golden-features), preserving voice. Verdict: **\*\*revise\*\***.

**\*\*Fix #2 ? revise.\*\*** The corrected\_text "the golden corpus + labelled videos live on the owner's machine" is imprecise: this IS the owner's machine. The accurate distinction is that the labelled golden data hasn't been GPU-extracted into the shared folder yet (features dir empty), and the heavier confirmation is owner-driven. I'll tighten to a claim I can fully stand behind: the GPU/torch/cv2 claim is removed (the real fix), and the owner-side framing kept minimal and accurate without overclaiming about machine location. Verdict: **\*\*revise\*\***.

Let me check the regression test referenced and confirm Gate A wiring framing isn't affected ? actually that's out of scope; both fixes are purely about environment-capability. I have enough to finalize.

### 15. user (2026-06-19T15:53:47.695Z)

Structured output provided successfully

### 16. assistant (2026-06-19T15:53:59.290Z)

Both proposed fixes survive, but I revised the corrected\_text on each because the verification surfaced a second-order inaccuracy in the proposals.

Summary of verification:

**\*\*Core staleness ? confirmed for both (file: `docs/NEXT\_STEPS.md`):\*\***

- Line 18: "the dev box is CPU-only: `torch+cpu`/`cv2` present but no GPU"
- Line 69: "Dev container no GPU/torch/cv2"

Both are false. On this machine (host `AVIS-MSI`, the owner's own "2070S" box the doc itself references at line 93): `torch 2.6.0+cu124` with `torch.cuda.is\_available() == True`, `cv2 4.13.0`, and `nvidia-smi` reports an RTX 2070 Super with Max-Q. Both stale\_text strings appear verbatim and uniquely.

**\*\*Why I revised (not confirmed) the proposed corrected\_text:\*\***

The proposals also re-asserted that the data is "not here" / "lives on the owner's machine." Verifying that clause showed it is partly wrong in the opposite direction:

- Videos ARE present: e.g. `C:\Users\avidu\Projects\Annotation Setup\Collect\Candidates\\*.MP4`
- Trajectory CSVs ARE present: `C:\Users\avidu\Projects\Annotation Setup\Collect\Trajectories\\*\_traj.csv`



- `output\human\_lovo\_manifest.json` exists locally
- This checkout IS the owner's machine, so "owner-side (GPU/WSL machine)" / "owner's machine" framing is imprecise.

What is genuinely absent: the `Golden Labelled` directory (the `.rallies.csv` labels ? directory does not exist) and `...\OneDrive\rally-annotator\golden-data\features\` (empty). I retargeted both corrected\_texts at those verifiably-missing artifacts so the edit fixes the GPU claim without introducing a new false "videos/CSVs not here" claim.

Output delivered via StructuredOutput (2 survivors, both verdict `revise`).