

claude-windows / 144a31b7-7962-4fee-9d6f-8a7854daa63f

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\144a31b7-7962-4fee-9d6f-8a7854daa63f\subagents\workflows\wf_1c5f1f5c-9ba\agent-adf6631ff8e5a5a2e.jsonl`
- SHA-256: `f0365eeca707d2f5e5f6e7b1688a820642ef18a3bf969f3111aa08deed2c1279`
- Source modified: `2026-06-20T08:01:32+00:00`
- Imported at: `2026-07-05T16:48:28+00:00`
- project: `wf_1c5f1f5c-9ba`
- session_id: `144a31b7-7962-4fee-9d6f-8a7854daa63f`
```

Transcript

1. user (2026-06-20T08:00:25.441Z)

Extract the EXACT structure a tracked design/project doc must follow in this project. Read: C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PROJECT_DOC_TEMPLATE.md (the canonical template), and skim docs/RALLY_QUALITY_RESEARCH.md and docs/PROMINENT_COURT_DETECTION.md as real examples of the template in use. Also read C:\Users\avidu\Projects\khelesutra\docs\REQUEST_ACCESS_FORM.md and C:\Users\avidu\Projects\khelesutra\docs\PORTABILITY.md to match the khelesutra repo's doc voice. Report: required_sections (in order), header_fields (the status-header fields like Status/Owner/Decisions-locked/etc.), the progress_tracker_format, the definition-of-done format, and style_notes (numbering convention, tone). Quote path:line. Return the schema.

2. assistant (2026-06-20T08:00:33.088Z)

I'll read the canonical template and the example docs, plus the khelesutra docs for voice matching.

3. user (2026-06-20T08:00:34.089Z)

```
1      <!--
2      PROJECT DOC TEMPLATE ? copy this file to docs/<YOUR_PLAN>.md and fill it in.
3
4      A "tracked project doc" is how this repo carries substantial design/project work: it can
5      be
6      ITERATED on, its progress TRACKED, and it is MARKED DONE against explicit deliverables,
7      then
8      archived. Standardizes the emergent pattern in docs/RALLY_QUALITY_RESEARCH.md §7 and
9      docs/GOLDEN_REGRESSION_FIXTURES.md.
10
11     Rules of thumb:
12     - §7 (the progress tracker) is the SOURCE OF TRUTH ? one row per independently-shippable
13     PR.
14     - Keep it HONEST: tag each claim [verified] (checked against code), [researched] (needs
15     sign-off),
16     or [design] (proposed). The doc reflects real shipped state, never aspirational.
17     - It POINTS to detail (code anchors, research artifacts), it does not duplicate it.
18     - Sections marked OPTIONAL are included only when relevant (e.g. Foundation/Threat-model
19     for
20     research- or security/privacy-heavy work). Delete sections that do not apply.
21     - Move to docs/archives/ once Status = DONE (per the archive policy: only terminal-state
22     work).
```

```

**Created:** `<YYYY-MM-DD>` . **Last updated:** `<YYYY-MM-DD>`
22 > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
23 > **Tracking anchors:** $7 progress tracker is the source of truth; indexed in
`docs/README.md`; pointer in memory `current-state` (and `docs/NEXT_STEPS.md` once work starts).
24 > **Relation to existing docs:** `<extends | supersedes | peer-of ?>`
25 > **Honesty note:** mark each claim `[verified]` / `[researched]` / `[design]`. Not
legal/financial advice where applicable.
26
27 ---
28
29 ## 0. TL;DR
30 <2?5 sentences: the goal, the approach, and the single most important caveat. Call out
any finding that is already actionable on its own.>
31
32 ## 1. Problem & goal
33 <Why now; the concrete user-facing outcome; what "good" looks like. Link the code/docs a
reader must read to get current.>
34
35 ## 2. Decisions locked
36 <Table ? capture answered questions so they are not relitigated.>
37
38 | # | Decision | Source / date | Implication |
39 |---|-----|-----|-----|
40 | D1 | ? | ? | ? |
41
42 ## 3. Foundation ? research / prior art *(OPTIONAL)*
43 <Include when the plan rests on external facts: cited research, legal/compliance,
benchmarks, prior-art survey. Mark researched-vs-verified; list any sign-off gates this
creates.>
44
45 ## 4. Design / architecture
46 <The actual plan. Data contracts, diagrams, and a REUSE MAP ? name the existing
modules/code this builds on so it is clear what is new vs reused.>
47
48 ## 5. Threat model / risk table *(OPTIONAL ? security / privacy / compliance work)*
49 <Who/what is protected vs NOT; residual risks stated plainly.>
50
51 ## 6. Honest limits ? what this does NOT do
52 <The false-confidence guardrail: what a green suite / a DONE status does NOT prove.>
53
54 ## 7. Deliverables & progress tracker ? **source of truth**
55 Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. **One small PR per row**,
branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.
56
57 | ID | Deliverable | Depends on | Gated? | Status | PR |
58 |---|-----|-----|-----|-----|----|
59 | P0 | ? | ? | No | ? | ? |
60
61 ## 8. Open questions ? owner / external *(split blocking-gates from nice-to-knows; name
who decides)*
62 1. ?
63
64 ## 9. Definition of done
65 <Explicit acceptance checklist. When all ticked ? flip `Status ? DONE` and archive.>
66 - [ ] ?
67
68 ## 10. References
69 **Internal code anchors** `[verified]`: **`<paths>`**. **Internal docs:** **`<docs>`**.
**External:** **`<cited sources>`**.
70
71 ---
72
73 ### Changelog
74 - `<YYYY-MM-DD>` ? `<what changed>`.

```

4. user (2026-06-20T08:00:35.502Z)

```

1      # Prominent-Court Detection ? kill background-court false positives (multicourt)
2
3      > **Status:** `DRAFT` ? design + experiment **owner-gated for any default flip** .
**Owner:** Avi . **Created:** 2026-06-19 . **Last updated:** 2026-06-19
4      > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
5      > **Tracking anchors:** $8 progress tracker is the source of truth; pointer in memory
`current-state`.
6      > **Relation to existing docs:** attacks gap **G2** ("precision ? local over-fires,
worst on multicourt") from [`RALLY_DETECTION_GAP_CLOSURE.md`](RALLY_DETECTION_GAP_CLOSURE.md);
it is the **shuttle-side** counterpart of the person-side **Gate B / court mask** in
[`MULTI_SIGNAL_FUSION_PLAN.md`](MULTI_SIGNAL_FUSION_PLAN.md). Complements the held-shuttle
**`shuttle_player_colocation`** (G3) ? court-localization (G2) and held-shuttle (G3) are
independent precision levers.
7      > **Honesty note:** claims are tagged `[verified]` (checked vs code/measured) /
`[design]` (proposed). Every signal lands **flag-gated, default-OFF**, promoted only on measured
lift per [`EXPERIMENT_HARNESS.md`](EXPERIMENT_HARNESS.md) + [[decision-rigor-norm]].
8
9      ---
10
11     ## 0. TL;DR
12     In **multicourt** footage the camera sees several courts. The shuttle detector (WASB)
tracks *any* shuttle in frame, so a rally on a **background / adjacent court** becomes a
predicted rally on the **main match** ? a false positive. The fix: identify the **prominent
court** (the foreground court that occupies the majority of the frame, where the match being
filmed is played) and **keep only rallies whose shuttle lives inside it**. Mechanically this re-
defines the shuttle signal from `(x, y, visible)` to `(x, y, visible-AND-inside-the-prominent-
court)` *before* windowing ? so a background-court shuttle reads as "not there" and never forms
a rally. A flat 2D floor polygon is the MVP; a **pseudo-3D play-volume** (court floor + the air
column above it) is the refinement so a high clear over the prominent court isn't wrongly
rejected. Lands as a **default-OFF experiment** ; the flip is gated on a Launch Report ([[launch-
report-convention]]).
13
14     ## 1. Problem & motivating example `[verified]`
15     - **Measured gap (G2):** local over-fires ~2x vs Gemini, and the over-fire is **worst on
multicourt** ? even among long (>7s) rallies, multicourt over-fires **1.58x** (background-court
games are *real* long rallies, just on the wrong court). See `RALLY_DETECTION_GAP_CLOSURE.md` $3
+ the P7 long-rally lens.
16     - **Concrete anchor ? GX010142 rally #1 (2026-06-19).** Owner observation: "the first
rally is just a shuttle flying in the non-prominent court? while the prominent-court players
have not started." Confirmed from the cached WASB trajectory
(`output/wasb/GX010142_1080p__wasb.csv`): rally #1 (`0.72?3.65s`, 2.94s) has a shuttle
centroid at **meanX ? 1616** on a 1920-wide frame (?84% to the right) ? spatially apart from the
central cluster where most rallies live (meanX ? 1000?1300). Several other suspicious rallies
(#3, #9, #12, #16) cluster in the same far-right band ? consistent with a **right-side adjacent
court**.
17     - **Why a 1D threshold is not enough `[verified]`:** the courts do not separate cleanly
on X or Y alone (rally Y-centroids overlap 250?620 px across courts; the far-right band mixes
with foreground play). The courts are **2D regions** under perspective ? so we need the court
*geometry*, not a coordinate cutoff. This is the owner's "pseudo-3D / court model" intuition,
validated by the data.
18
19     ## 2. Goal & non-goals
20     - **Goal:** on multicourt clips, drop rallies whose shuttle is outside the prominent
court, **without losing real prominent-court rallies** ? measured per-video, no regression on
single-court clips (where it must be a no-op or a do-no-harm).
21     - **Non-goals (v1):** multi-camera; tracking players across courts; full court-line
homography for scoring; anything biometric. Single-court footage is explicitly a **no-op** (one
court ? everything is "prominent").
22
23     ## 3. Decisions locked

```

```

24      | # | Decision | Rationale |
25      |---|-----|-----|
26      | D1 | The prominent court is a **2D region (polygon) in image space**, refined to a
**pseudo-3D play volume** | 1D thresholds don't separate courts under perspective (§1) |
27      | D2 | Gate the **shuttle trajectory** by the region **before windowing** (re-define
visibility) | Stops a background rally at the source ? no window ever forms there. Reuses
`point_in_polygon` |
28      | D3 | **Default-OFF, flag-gated; single-court = no-op** | Zero-regression rollout
([[weak-signals-retained]]); no behaviour change unless a polygon is supplied + the flag set |
29      | D4 | Reuse the existing **`court_polygon`** config + `point_in_polygon`; do NOT invent
a new geometry stack | `FusionConfig.court_polygon` + `fusion_features.point_in_polygon` already
exist `[verified]` |
30      | D5 | The prominent court can be **supplied (config) OR auto-derived**; auto-derivation
is its own measured step | Config polygon is the cheapest first experiment; auto is the product
path |
31
32      ## 4. Design
33
34      ### 4.1 The signal change (the core experiment)
35      Today the windower consumes the raw WASB trajectory: a per-frame `(Frame, Visibility, X,
Y)`. The change is a **pre-windowing filter**:
36
37      ```
38      for each frame f: if shuttle_visible(f) and NOT inside_prominent_court(X_f, Y_f): mark
f as NOT visible
39      ```
40
41      A background-court shuttle is then indistinguishable from "no shuttle" ? the velocity
windower never opens a rally there. This is purely **subtractive** on the shuttle stream (it can
only remove visibility, never add it), so on a correctly-drawn prominent court it cannot
manufacture rallies ? only suppress out-of-court ones. (Mirror of `rally_gate`'s "do-no-harm"
discipline.)
42
43      ### 4.2 Identifying the prominent court (four routes, cheapest first)
44      1. **Configured polygon** (MVP, `[design]?buildable now`): a normalized `[[x,y],?]'`
court polygon per camera/clip, exactly the existing `FusionConfig.court_polygon`. Draw it once
from the dense shuttle/player cluster. Good enough to validate the whole concept on GX010142 +
the multicourt golden clips.
45      2. **Player-location auto-derivation** (`[design]`): run
`PersonDetector.detections_per_frame` `[verified]` exists, take the **largest, most-central,
most-persistent** player cluster (the foreground players are bigger boxes lower in frame), and
fit the prominent-court polygon to where *those* players' feet live. "Prominent" = where the
match's players are." This is why the owner asked to **extract person features** ? players
anchor the court.
46      3. **Shuttle-density auto-derivation** (`[design]`): the prominent court is the densest
sustained shuttle-activity region; cluster the visible trajectory and take the dominant mode.
Person-free fallback.
47      4. **Court-line detection** (`[design]`, heaviest): Hough/segmentation court lines ? the
largest court quad. Most accurate, most work; deferred.
48
49      v1 ships route 1 (config) + measures; route 2 (player-anchored) is the product upgrade.
50
51      ### 4.3 Pseudo-3D play volume (the refinement)
52      A flat floor polygon rejects a shuttle at the top of a high clear *over the prominent
court* (it projects above the floor quad). The pseudo-3D model treats the prominent court as a
**floor quad + a vertical air column**: accept a shuttle if its image point lies within the
floor polygon **expanded upward** by a height envelope (the max plausible shuttle height
projected through the camera). Practically: the floor quad plus an upward `pad` (perspective-
aware ? larger near the net/far baseline). Start with a generous fixed upward pad (cheap, MVP);
upgrade to a per-camera height envelope only if measurement shows real rallies being clipped.
This keeps the floor tight (rejects adjacent-court floor play) while not clipping legitimate
high shuttles.
53
54      ### 4.4 Reuse map `[verified]`
55      | Need | Existing code |

```

```

56     |---|---|
57     | point-in-region test |
`backend/pipeline/detectors/fusion_features.py::point_in_polygon` |
58     | normalized court polygon config |
`backend/config/models.py::FusionConfig.court_polygon` (today: person/Gate-B only) |
59     | player boxes per frame |
`backend/pipeline/detectors/person_detector.py::PersonDetector.detections_per_frame` |
60     | trajectory parse + windowing |
`backend/pipeline/detectors/tracknet_runner.parse_trajectory_csv`, `backend/eval/windowing.py` |
61     | per-video scoring | `backend/eval/rally_seg_eval.py` (+ the `--stratify` multicourt
split, #220) |
62
63     New code is small: a `prominent_court_gate(points, polygon, height_pad)` that filters
trajectory points, threaded into the windowing input behind a default-OFF flag.
64
65     ### 4.5 Extension ? deterministic rally-END via shuttle-floor contact (owner note,
2026-06-19) `[design]`
66     The same court geometry that localizes the prominent court also unlocks a
**deterministic rally-END signal** ? the cheapest, most certain way to know a point is over:
**the shuttle hit the floor.** This attacks **G3** (rally-END over-extension), where today's
shuttle-velocity windowing runs the window long because "shuttle still moving ? rally over."
67
68     - **Introduce two geometric primitives:**
69     - **The white court lines.** A badminton court is bounded by **white, straight lines
forming a rectangle** (with the perspective trapezoid in image space). Detect them ? Hough lines
/ line-segment detection filtered to white, long, court-consistent segments ? to recover the
**court quad** (which also is the prominent-court floor polygon of §4.2 route 4; one detector
serves both).
70     - **The floor plane.** The court lines define the **floor** (the plane the players
stand on). The shuttle in flight lives above this plane; a shuttle at the plane (within the
trapezoid, or just outside it) is on the floor.
71     - **The rally-END rule (deterministic):** a rally ends when the shuttle makes floor
contact ? its trajectory descends to and stops/rests at the floor level (a near-zero-
velocity terminus at the court-line plane), OR a shuttle that goes to the floor outside the
prominent court = an out-of-bounds landing (also point-over). Unlike velocity heuristics, a
shuttle resting on the floor is unambiguous ? the point is over.
72     - **Why it's strong + cheap:** no new model, no person stack ? it reuses the WASB
trajectory (which already gives the shuttle's descending `(x, y)` + the `Visibility` drop when
it settles) plus the court-line geometry. It is the shuttle-side rally-END signal,
complementary to the player-state rally-END (player kinematics / held-shuttle) in
`RALLY_DETECTION_GAP_CLOSURE.md` §4 Lever A ? fuse both for robustness (a shuttle can settle
off-camera; a player can pick up early).
73     - Subtlety to measure: distinguish a true floor landing (point over) from a shuttle
that dips low but is played back up (a tight net shot / low retrieve). The court-line plane
+ a brief dwelling at floor level (the shuttle rests, doesn't re-accelerate up) is the
discriminator; tune the dwell window on golden clips.
74
75     This is tracked as C8 below and cross-links gap-closure G3. It shares the court-
line detector with the prominent-court polygon (§4.2 route 4), so the two land together once
court-line detection exists.
76
77     ## 5. Experiment plan (measurement)
78     Land the gate flag-gated, then measure on the cached trajectories (CPU, no GPU/re-
detect):
79     1. GX010142 sanity: draw a prominent-court polygon; confirm rally #1 disappears
and the central rallies survive. The litmus the whole idea is born from.
80     2. Multicourt golden (`mahadevpura_2`, GX010128, Badminton_BXH_2, Boxhill_Doubles`):
? rally-count, and once those clips are golden-labelled, ? precision / ? recall vs golden via
`rally_seg_eval --stratify` (single vs multicourt). The win condition: multicourt precision up,
recall flat (we drop FPs, not real rallies).
81     3. Single-court no-op: on single-court clips with no polygon (or a full-frame
polygon) the gate is a verified no-op ? assert zero change.
82     4. Dual-eval discipline ([[eval-dual-set-regression]]): no per-video regression on
either eval set.
83

```

```

84     Promotion to default-ON is a **separate, owner-gated** decision with a Launch Report
(both rulers + the over-seg guard), never bundled with the experiment PR.
85
86     ## 6. Risks & limits
87     - **Polygon is camera-dependent.** A hand-drawn polygon doesn't generalize across setups
? route 2/3 (auto-derivation) is the real product path; v1 config-polygon is for *validation*.
88     - **Shuttle height** ($4.3): too-tight a floor clips high clears; the pseudo-3D pad
mitigates, measurement tunes it.
89     - **The prominent court can change** (camera re-aim mid-clip): out of scope v1 (assume
static per clip).
90     - **Not a recall fix.** This only removes out-of-court FPs; it does nothing for missed
in-court rallies (G1) ? and must be proven not to *cost* recall.
91
92     ## 7. Open questions (owner / data)
93     1. Auto-derivation anchor: **player-location** (route 2) vs **shuttle-density** (route
3) as the default once validated? *(decides: data, post-experiment)*
94     2. Height-envelope: fixed upward pad vs per-camera perspective model ? is the fixed pad
enough? *(decides: measurement)*
95     3. Do we expose the prominent court as a **one-time per-camera setup** (draw-once UI) or
always auto? *(decides: owner/product)*
96
97     ## 8. Deliverables & progress tracker ? source of truth
98     Legend: ? Todo · ? In progress · ? Done · ? Gated. One small PR per row, off
`origin/master`, no stacking.
99
100    | ID | Deliverable | Depends on | Gated? | Status | PR |
101    |----|-----|-----|-----|-----|----|
102    | C0 | This design doc | ? | No | ? | this PR |
103    | C1 | `prominent_court_gate(points, polygon, height_pad)` + unit tests (pure, point-in-
polygon) | ? | No (default-OFF) | ? | ? |
104    | C2 | Thread the gate into the windowing input behind a default-OFF flag
(`court_polygon` + `gate_shuttle_to_court`) | C1 | No (default-OFF) | ? | ? |
105    | C3 | GX010142 litmus: rally #1 dropped, central rallies kept (measured) | C2 | No | ?
| ? |
106    | C4 | Extract person features (`detections_per_frame`) on GX010142 ? player-cluster =
prominent court; measure `shuttle_player_colocation` too | ? | No | ? | ? |
107    | C5 | Multicourt-golden measurement (? count now; ? precision/recall vs golden once
labelled) | C2, labels | ? owner labels | ? | ? |
108    | C6 | Auto-derive the polygon (route 2 player-anchored) | C4 | ? owner | ? | ? |
109    | C7 | Launch decision (default-ON) with a Launch Report | C3,C5 | ? owner | ? | ? |
110    | C8 | **Court-line / floor detection** ? (a) auto court polygon, (b) **deterministic
shuttle-floor-contact rally-END** ($4.5, targets G3) | ? | ? owner | ? | ? |
111
112     ## 9. Definition of done
113     - [ ] The gate exists, is **default-OFF**, unit-tested, and a verified no-op on single-
court / no-polygon.
114     - [ ] GX010142 rally #1 (and the far-right adjacent-court band) is dropped while the
central rallies survive `[measured]`.
115     - [ ] On the multicourt golden clips, multicourt **precision improves with recall flat**
vs golden labels (no per-video regression).
116     - [ ] A winning configuration is flipped default-ON only via a Launch Report ? flip
`Status ? DONE`, archive.
117
118     ## 10. References
119     **Code `[verified]`: **`fusion_features.py` (`point_in_polygon`, `person_in_court`),
`person_detector.py` (`detections_per_frame`), `config/models.py`
(`FusionConfig.court_polygon`), `eval/windowing.py`, `eval/rally_seg_eval.py` (`--stratify`).
120     **Docs: **[`RALLY_DETECTION_GAP_CLOSURE.md`](RALLY_DETECTION_GAP_CLOSURE.md) (G2),
[**MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (Gate B / court mask),
[**EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md),
[**QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md).
121
122     ---
123
124     ### Changelog

```

125 - 2026-06-19 ? Created from the owner's GX010142 rally-#1 (background-court FP)
observation + the measured multicourt over-fire (G2). Validated the 2D/pseudo-3D need from the
cached trajectory; mapped the reuse onto the existing `court\_polygon` + `point\_in\_polygon`
machinery.
126 - 2026-06-19 ? Added §4.5 (owner note): a **deterministic rally-END signal via shuttle-
floor contact** ? detect the white court lines ? the floor plane ? a shuttle resting on the
floor (or landing out-of-court) ends the rally (targets G3). Shares the court-line detector with
the prominent-court polygon; tracked as C8.
127

5. user (2026-06-20T08:00:35.523Z)

```
1      # Rally-Detection Quality: Research, Eval Framework & Project Plan
2
3      **Status:** Foundation doc (2026-06-16), run as a **tracked project** (§7) ? **"Bridge
the
4      local?Gemini rally-detection gap,"* **completed when all committed tiers land and a
single quality
5      number is attributed (§7.1).** The empirical + research backbone for closing the gap
between the
6      **local, no-cloud rally segmenter** and the **Gemini** path. Iterations land as their
own measured
7      PRs, logged in [`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md); this doc is the plan
they execute against.
8
9      **How to read this:** §1 frames the gap with measured numbers. §2 is the precise, code-
grounded
10     diagnosis (read before proposing fixes). §3 inventories the eval/experiment harness we
already
11     have. §4 specifies how we strengthen that harness (the "really strong evaluation
framework").
12     §5 is the cited external-research synthesis (4 dimensions). §6 is the prioritized
roadmap with
13     falsifiable success criteria. **§7 is the project plan ? work-item tracker, sequencing,
the single
14     number we attribute to the effort, and the definition of done.** §8 recaps the non-
negotiable
15     constraints + cross-links.
16
17     > **This doc does NOT duplicate prior design.** It builds on, and cross-links:
18     > [`HOW_RALLY_DETECTION_WORKS.md`](HOW_RALLY_DETECTION_WORKS.md) (the 2-layer mental
model),
19     > [`EXPERIMENT_HARNESS.md`](EXPERIMENT_HARNESS.md) (the A/B + LOVO harness, P1 shipped),
20     > [`MULTI_SIGNAL_FUSION_PLAN.md`](MULTI_SIGNAL_FUSION_PLAN.md) (shuttle+person+audio
fusion),
21     > [`ANALYZERS_AND_RECONCILERS.md`](ANALYZERS_AND_RECONCILERS.md) (the signal-fusion
framework +
22     > §13 literature review ? **already** establishes that our skeleton is standard TAL/TAS
+
23     > late-fusion; we *cite, don't re-derive*), [`ABLATION_LAYER.md`](ABLATION_LAYER.md),
24     > [`EVALUATION.md`](EVALUATION.md),
[`OWNED_MODEL_IMPLEMENTATION_PLAN.md`](OWNED_MODEL_IMPLEMENTATION_PLAN.md)
25     > (TAL-head-first ? VLM), and [`GOLDEN_SET_VENDORING.md`](GOLDEN_SET_VENDORING.md)
(corpus growth).
26
27     ---
28
29     ## 1. The gap (measured)
30
31     Two paths produce rally segments today. **Gemini** (cloud VLM, watches the video
semantically)
32     is strong; the **local** path (WASB shuttle trajectory ? velocity windowing, no cloud)
lags
33     badly. Measured F1 vs golden labels ([`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md)):
34
```

```

35 | Path | Mahadevpura (clean, single court) | Kushagra (multi-court) |
36 |---|---|---|
37 | Heuristic velocity windowing (local today) | 0.657 | **0.157** |
38 | Gen-0 owned head (LogReg on 5 trajectory features, LOVO n=6) | 0.744 | 0.561 |
39 | Two-stage WASB + Gemini refine | 0.83 | ? |
40 | **Gemini wrapper (gemini-flash)** | **0.93** | **0.66** |
41
42 The local heuristic is **footage-fragile**: fine on a clean single court, collapses on
43 multi-court / busy venues. The owned Gen-0 head already closes most of the multi-court
gap
44 (0.561 vs 0.66) ? the residual is **corpus-sized, not architecture-sized**.
45
46 ### 1.1 The 2026-06-16 local-vs-Gemini divergence run (no golden labels)
47
48 A fully-local run (`wasb_hybrid --skip-ai-handoff`) vs Gemini on 3 fresh matches
49 (Kushagra/Yash June-12; report: `~/June 12 Local-vs-Gemini/_LOCAL_vs_GEMINI_REPORT.md`,
50 generator `scratch/compare_local_vs_gemini.py`). **Gemini is a strong reference here,
NOT
51 ground truth** (these clips are unlabeled) ? so this is an agreement/divergence study:
52
53 | | Gemini | Local (WASB no-AI) |
54 |---|---|---|
55 | rallies / windows (3 videos) | 59 | 12 |
56 | total "play" | 6:28 | **17:41** |
57 | mean window duration | **6.1 s** | **65 s** (10.6x) |
58 | true misses (gemini-only) | ? | **0** |
59 | merge groups (one local window ? 2 Gemini rallies) | ? | 7 |
60
61 Headline: **the local path does not *miss* rallies ? it fails to *segment* them.** One
62 mahadevpura-1 window spanned the entire 174 s clip. Local "play" is 2.7x Gemini's
because its
63 mega-windows swallow the dead time between rallies.
64
65 ---
66
67 ## 2. Diagnosis (code-grounded ? read before proposing fixes)
68
69 The local pipeline is **two layers** (see
[HOW_RALLY_DETECTION_WORKS.md])(HOW_RALLY_DETECTION_WORKS.md)):
70 ***(L1) detection** ? per-frame shuttle position (WASB/TrackNet); and ***(L2)
segmentation** ?
71 turn the trajectory into rally intervals. The gap is almost entirely **L2**.
72
73 ### 2.1 L1 (shuttle tracking) is NOT the bottleneck
74 On the divergence clips WASB reported the shuttle visible in **96?99 % of frames**
(mahadevpura-1
75 98.6 %, GX020125 72.8 %, mahadevpura-2 78.8 %). The trajectory signal is present and
dense.
76
77 ### 2.2 L2 (segmentation) is the bottleneck ? and the flaw is conceptual
78 The windowing is `trajectory_to_action_windows`
79 ([`backend/pipeline/detectors/tracknet_runner.py:256`](../backend/pipeline/detectors/tra
cknet_runner.py)):
80
81 - A frame is **"active"** iff the shuttle is **visible** AND its **per-second-normalized
speed**
82 exceeds `velocity_threshold` (default 0.02). Velocity = `(dist / frame_width) * (fps /
dframes)`
83 ? **already fps-normalized** (? fraction of frame-width per second).
84 - Active frames within `dead_time_merge_gap` (2.0 s) are **merged** into one window.
85 - Windows shorter than `min_action_window_duration` (2.0 s) are dropped.
86
87 **The conceptual flaw: "the shuttle is moving" is treated as "a rally is in progress."**
During
88 dead time the shuttle is still visible and gently moving (being picked up, walked back,

```


knock-ups,

89 slow serve prep) ? those frames read as **active**. The 2 s merge then bridges the brief gaps, and

90 because (at the time this was diagnosed) there was ***no `max\_window\_duration` cap and no*

91 *inactivity/boundary split***, a single window grew until the shuttle finally went invisible ? which,

92 at 96?99 % visibility, is rarely. Result: rally + dead-time + rally + ? collapse into a few

93 mega-windows. **(This is the failure W1/R4 fix; both are now default-ON since the R1 milestone, PR*

94 *#204 ? `max\_window\_duration=6`, `window\_inactivity\_end=1.5` in `IndexingConfig`.)**

95

96 > ***Correction to an earlier hypothesis (honesty note).*** A first pass guessed the cause was

97 > "velocity threshold not fps-invariant." It IS fps-invariant (the *`fps/dframes`* factor). The

98 > real cause is the **motion-equals-rally** conflation + 2 s bridging + no upper bound. fps matters

99 > only secondarily: denser 60 fps sampling keeps more slow-motion dead-time frames above the

100 > speed threshold. The *`\_LOCAL\_vs\_GEMINI\_REPORT.md`* has been corrected to match.

101

102 *### 2.3 What this implies*

103 The local path ***has no rally-STATE signal*** ? only "is the shuttle moving." Bridging the gap is

104 fundamentally about giving L2 signals that separate **rally** from **shuttle-merely-moving**

105 (net-crossings, two-sided play, sustained fast exchange, hit sounds) and a segmenter that

106 produces **bounded, well-cut** intervals ? exactly the multi-signal-fusion + learned-segmenter

107 direction the existing docs lay out. This doc makes that concrete and measurable.

108

109 ---

110

111 *## 3. What we already have (eval/experiment harness inventory)*

112

113 The harness is rich and composable ? we ***extend*** it, not reinvent it. (Modules under

114 *`backend/eval/`*.)

115

116 ***Scoring & matching***

117 - *`metrics.py`* ? temporal-IoU interval matching; precision/recall/F1, mean_iou, boundary errors.

118 - *`segmentation\_metrics.py`* ? ***F1@{0.1,0.25,0.5}***, ***mAP@tIoU*** (0.3?0.7), ***segment_count_ratio***

119 (the over/under-segmentation diagnostic). (Edit/segmental score deliberately omitted to date.)

120

121 ***Statistical rigor (small-n)***

122 - *`eval\_stats.py`* ? bootstrap 95 % CIs, paired exact ***sign-test***, *`paired\_delta\_ci`*,

123 *`grouped\_folds`* (LOGO), *`out\_of\_fold\_predictions`* (anti-stacking-leakage). The promotion bar =

124 "F1 CI excludes 0 ***AND*** sign-test agrees."

125 - *`score\_calibration.py`* ? Platt + isotonic calibrators, ***Brier + ECE***.

126

127 ***Variants, A/B, leaderboard***

128 - *`experiment.py`* ? *`Variant`* (declarative recipe: segmenter + config + cue_flags + gates +

129 learned-classifier + threshold/merge_gap/calibrate), *`compare()`* (LOVO-honest A/B with an

130 "honest" stats block), *`compare\_unlabeled()`* (agreement on unlabeled footage), *`CONTROL`*

131 (pinned baseline), *`record\_experiment()`* ? *`eval\_baselines/experiment\_leaderboard.json`*.

132 - *`ablation.py`* ? leave-one-signal-out marginal ?F1 with CI + sign-test + a no-over-claim verdict

```

133     vocabulary (`earns_keep` / `suggestive` / `inert` / `structural_protected`).
134     - `classifier.py` ? dependency-free numpy `LogisticRegression` (the learned scorer;
JSON-serializable).
135
136     **Eval?serve discipline + gates**
137     - `windowing.py` ? single-sources the **SERVED** preset from `IndexingConfig` defaults
so eval
138     and production can't silently diverge; `PROPOSAL` (recall-oriented) is where the
golden feature
139     substrate is built. (This exists *because* Gate-A measured **+0.074 on PROPOSAL but
?0.008 on
140     SERVED** and was reverted ? the canonical eval?serve lesson.)
141     - `promotion.py` / `served_gate_a.py` ? promote a cue only if it does NOT regress at the
served
142     windowing; `per_user_gate.py` / `per_user_regression.py` ? per-user analogues;
`regression.py` ?
143     golden-set no-regression baseline with a GT-hash to detect label changes.
144
145     **Golden substrate**
146     - `gt_loader.py` ? one canonical GT loader (both legacy formats). `fusion_golden.py` /
147     `fusion_compare.py` / `fusion_audio.py` ? extract replayable per-window feature JSONs
148     (shuttle 5 + person 5 + audio 3) so ablation/compare re-score on CPU without re-
detecting.
149     - **Corpus today: 15 golden videos** (started at the 6-video label-set hash
`50d98ffbc370`, since
150     grown with the GX/Boxhill multicourt clips); more expected via the
151     Roora vendoring track ([`GOLDEN_SET_VENDORING.md`](GOLDEN_SET_VENDORING.md)).
152
153     **Signals available to L2 (built, default-OFF):** net-crossing count (`rally_gate.py`,
Gate-A),
154     5 person/court features (`person_detector.py` + `fusion_features.py`), 3 audio-hit
features
155     (`audio_features.py`). Fusion hooks: `fusion_hybrid.py::_enrich_candidate_stats` /
156     `_select_for_handoff`. Persisted in `candidate_segments.stats`.
157
158     ### 3.1 Known gaps in the harness (what $4 fills)
159     1. **Over-segmentation is under-surfaced in the default report.** `segment_count_ratio`
exists but
160     the merge/split *breakdown* (the very thing that exposed the over-merge) lives only
in a scratch
161     script. tIoU-F1 alone is **blind to over-merge** ? a giant window can still "match"
via overlap.
162     2. **No segmental edit-score** (Levenshtein over the predicted/gt label sequence) ? the
standard
163     TAS over-segmentation metric.
164     3. **Per-cue calibration is fitted but never *measured* post-hoc** (no ECE/Brier delta
reported).
165     4. **`out_of_fold_predictions` (anti-leakage stacking) is built but not wired to an
arbiter.**
166     5. **No active-learning loop** ? nothing tells us *which* unlabeled video to label next.
167     6. **No single "rally-segmentation eval" entrypoint** ? the pieces exist but a one-
command
168     `score this segmenter vs golden, honestly` runner would make every iteration cheap.
169     7. **The local-vs-Gemini agreement track is ad-hoc** (a scratch script), not a first-
class
170     *shadow* signal alongside the golden gate.
171
172     ---
173
174     ## 4. The strengthened evaluation framework (the "really strong harness")
175
176     Design principles (inherited, non-negotiable):
177     - **Signals ? learned scorer, NO special-casing**
([`ANALYZERS_AND_RECONCILERS.md`](ANALYZERS_AND_RECONCILERS.md) $1):
178     every new cue is a feature column the scorer learns to weight; never an `if/else` in

```

the decision path.

```

179 - **Promote only on measured lift; default OFF; revert is one flag**
([`EXPERIMENT_HARNESS.md`](EXPERIMENT_HARNESS.md) §6).
180 - **eval ? serve:** a harness win must be re-confirmed on the *served* windowing before
any default flip.
181 - **Honest small-n stats:** never a bare mean at  $n \geq 16$ ; always CI + paired sign-test
(C1).
182
183 ### 4.1 Metric set the leaderboard should standardize on
184 Surface ALL of these in the default `compare()` report (most already computed; the task
is to
185 *report* them together so over-merge can never hide again):
186
187 - **tIoU-F1 @0.5** (primary) **and F1@{0.1,0.25,0.5}** (segmental, TAS-standard).
188 - **mAP@tIoU** (0.3?0.7) when predictions carry confidences.
189 - **`segment_count_ratio`** + a **merge-rate / split-rate** breakdown (fraction of GT
rallies that
190 share a predicted window = under-segmentation; fraction of predicted windows that span
?2 GT =
191 the over-merge signal). Promote the scratch `compare_local_vs_gemini.py` matching into
a reusable
192 `backend/eval/` helper.
193 - **Segmental edit (Levenshtein) score** ? the canonical over-segmentation metric we
currently lack.
194 - **Boundary timing** (mean start/end error on matched pairs) ? already in `metrics.py`,
surface it.
195 - **Calibration (ECE/Brier)** reported pre/post per-cue calibration (close gap #3).
196
197 > **Why:** plain tIoU-F1 rewarded our mega-windows (they "overlap" real rallies). The
metric set
198 > above makes over/under-segmentation *first-class*, so a fix that improves segmentation
is
199 > visible even when tIoU-F1 barely moves ? and vice-versa.
200
201 ### 4.2 Splits & significance
202 - **Leave-one-video-out, grouped by match** (never leave-one-*window*-out ? windows
within a video
203 are correlated). Keep `grouped_folds`. Note the **RL00CV bias** caveat for tiny n (see
§5.4).
204 - **Bootstrap CI + paired exact sign-test** on per-video held-out F1 (C1). A result is
*promotable*
205 only if CI excludes 0 **and** the sign-test agrees, **and** it survives the served-
windowing check.
206 - **Anti-leakage:** per-LOVO-fold threshold tuning + calibration on the TRAIN fold only
(already
207 enforced in `experiment._calibrate_and_tune`); wire `out_of_fold_predictions` once an
arbiter exists.
208
209 ### 4.3 The golden corpus + active learning (the highest-leverage lever)
210 The single biggest quality lever is **more labeled matches**
([`NEXT_STEPS.md`](NEXT_STEPS.md):
211 "corpus growth = highest leverage"). With ~6 ? ~16 videos:
212 - Grow via the Roora vendoring flow
([`GOLDEN_SET_VENDORING.md`](GOLDEN_SET_VENDORING.md),
213 [ `ROORA_LABELING_GUIDELINES.md`](ROORA_LABELING_GUIDELINES.md)) ? court calibration,
IAA, the
214 3 non-negotiables (ignore dead time, every rally, one-contact = one shot).
215 - **Active learning:** stop labeling videos at random. Score the *unlabeled* corpus and
label the
216 videos where the local segmenter is most *uncertain* / most *divergent from Gemini* /
most
217 *diverse* (new venue, fps, lighting). §5.3 surveys the methods; the local-vs-Gemini
agreement
218 signal (§4.4) is a ready proxy for "where do we disagree most ? label that."
219

```

220 ### 4.4 Shadow evaluation: Gemini-agreement as a free, unlimited signal (NOT a gate, NOT training data)

221 We have exactly one strong reference available on *unlabeled* footage at inference time: **Gemini**.

222 - `experiment.compare\_unlabeled()` already scores agreement (agreed / a-only / b-only) without labels.

223 - Use it as a **shadow metric**: run any candidate segmenter on a large unlabeled pool, measure

224 agreement-with-Gemini, and watch it as a *trend* ? cheap, unlimited, catches regressions the 6

225 golden videos can't. **It is a shadow signal only.** Golden labels remain the gate; Gemini is not

226 ground truth (it has its own errors) and ? per the licensing guardrail (§8) ? **its outputs may

227 never become training data for owned weights.**

228

229 ### 4.5 One-command rally-segmentation eval (close gap #6)

230 Wrap the above into a single reusable endpoint (a future PR ? not built here): given a segmenter

231 variant + the golden manifest, emit the full §4.1 metric set + the honest stats block + a leaderboard

232 row, and (optionally) the §4.4 shadow agreement on an unlabeled pool. Every iteration in §6 is then

233 "run the endpoint, read the honest delta, promote-or-not."

234

235 ---

236

237 ## 5. External research synthesis (cited, adversarially verified)

238

239 > Produced by a deep-research run (6 search angles ? 27 sources ? 126 candidate claims ? 25

240 > 3-vote-adversarially-verified, 23 confirmed / 2 refuted). **Cite, don't over-claim.** Two claims

241 > were *killed* in verification ? recorded in §5.5 so we don't repeat them. **License caveat:** the

242 > code repos below are research releases; **independently confirm each repo's actual license is

243 > MIT/Apache-2.0/BSD before vendoring any of it** (this survey did not verify every license file).

244 > **Governance:** none of these techniques require training on paid-LLM outputs ? all are compatible

245 > with our no-Gemini-training-data guardrail (§8).

246

247 ### 5.1 Windowing & boundaries (TAL/TAS) ? the explicit-boundary machinery we lack

248 The recurring lesson: modern action-localization models have an **explicit boundary mechanism**;

249 our heuristic has none (it merges on a velocity+gap rule). The transferable *idea* is "predict a

250 boundary and cut there," but the deep models are built for **hundreds+ of training videos**, so at

251 n=6?16 they are **research bets**, not drop-ins.

252

253 - **ASRF ? Action Segment Refinement Framework** (Ishikawa et al., WACV 2021, [arXiv:2007.06866](https://arxiv.org/abs/2007.06866)).

254 A shared feature extractor feeds **two branches**: frame-wise classification + a **Boundary

255 Regression Branch**; predicted boundaries refine/split the segmentation (+13.7% edit, +16.1%

256 segmental F1 over MS-TCN-era SOTA on 50Salads/GTEA/Breakfast). **Why it fits:** the boundary head

257 is exactly the missing "cut here" mechanism. **? Two caveats:** (1) ASRF's headline goal is to

258 *reduce* over-segmentation (merge spurious fragments) ? the opposite of our over-*merge*; only the

259 *split-at-a-detected-boundary* sub-mechanism transfers. (2) The claim that ASRF's

boundary branch

260 is a **model-independent post-hoc splitter** you can bolt onto any dense signal was
REFUTED

261 (\$5.5) ? you must train the two-branch model, not graft the branch. **Effort/payoff:**
research bet.

262

263 - **TriDet** (Shi et al., CVPR 2023,
[arXiv:2303.07347](https://arxiv.org/abs/2303.07347)). A

264 **Trident-head** models each boundary as a *relative probability distribution over
local bins*

265 (boundary = expected value over neighboring instants) ? built for imprecise/ambiguous
boundaries.

266 69.3% avg mAP on THUMOS14 (> ActionFormer 66.8% at ~75% latency), **convolution-only
(SGP), no**

267 **self-attention** (attractive on a single local GPU). Ablation shows the Trident-head's
gain

268 concentrates at strict tIoU=0.7 ? it buys *boundary precision*. **Effort/payoff:**
research bet

269 (trained on 100s of videos); the portable piece is the relative-boundary idea.

270

271 - **ActionFormer** (ECCV 2022, [arXiv:2202.07925](https://arxiv.org/abs/2202.07925)).
The reference

272 single-stage anchor-free template: multiscale pyramid + local self-attention; classify
every

273 moment + regress its boundaries (71.0% mAP@0.5 THUMOS14). This is the family
TriDet/TemporalMaxer/

274 CausalTAD extend, and the natural target if we ever graduate from heuristic+LogReg.
Research bet.

275

276 - **TemporalMaxer** (CVPRW 2023, [arXiv:2303.09055](https://arxiv.org/abs/2303.09055)) ?
the most

277 important small-data signal in this dimension. It replaces ActionFormer's self-
attention block

278 with a **parameter-free max-pool** and *beats* it (67.7 vs 66.8 avg mAP) at **~4x
fewer params**.

279 The ablation is the takeaway: a **learned 1-D conv block is the WORST** (59.4,
overfits redundant

280 features); parameter-free pooling wins. **Why it fits:** direct evidence to **keep our
scorer**

281 **parameter-light** (the numpy LogisticRegression over a handful of features is the
right choice at

282 our scale, not a placeholder to rush past). **Effort/payoff:** a discipline (free) ?
adopt the

283 principle, not the model. (Caveat: the overfitting evidence is about feature
redundancy on a

284 mid-size set, a directional inference to our small-n setting.)*

285

286 - **CausalTAD** (2024, [arXiv:2407.17792](https://arxiv.org/abs/2407.17792)) ?
steering-away note.

287 Causal-attention masking (attend only to past/present) won EPIC-Kitchens/Ego4D 2024
tracks, but it

288 adds **no** boundary-split module. Useful later for *online/streaming* rally
detection; **NOT** the

289 over-merge fix. Listed so we don't mistake causal attention for the remedy.

290

291 - **Metrics** (the actionable cross-cut): the field leads with **segmental Edit
(Levenshtein) score**

292 and F1@k, not frame accuracy (MoF) or plain tIoU ? see §5.4.

293

294 **### 5.2 Shuttle-trajectory ? rally signal processing ? published badminton systems
already do what we lack**

295 This is the highest-confidence, lowest-effort dimension: in-domain papers operationalize
the exact

296 boundary heuristics we're missing, in spirit license-compatible (re-confirm before
vendering code).

297

298 - ****See et al. ? Service & End-of-Rally Detection in Badminton**** (2024 Springer chapter

299

[10.1007/978-3-031-69073-0_15](https://link.springer.com/chapter/10.1007/978-3-031-69073-0_15);

300 2025 SN Computer Science follow-on

[10.1007/s42979-025-03935-0](https://link.springer.com/article/10.1007/s42979-025-03935-0)).

301 ******(a) Rally END = a sustained-inactivity rule:****** the rally ends when ******?80% of

consecutive frames

302 show no shuttle motion****** (~******95% accuracy****** over 187 rallies). ******(b) Rally START

(service) = a

303 learned classifier****** over court geometry + player poses (~26 features/frame × 12

frames ?

304 CNN/BiLSTM, ******88.6%****** accuracy). ****Why it fits:****** the inactivity rule is **precisely****

the

305 end-of-rally cut our windowing lacks ? a few lines, directly breaks 65 s mega-windows.

****? Caveats:********

306 the rule is a **termination**, not a **within-rally split**; results rest on paywalled

abstracts +

307 small test sets (35?187 rallies); the serve classifier needs the (default-OFF) person

cue + court

308 polygons. ****Effort/payoff:****** the inactivity rule = ******low-effort high-confidence win******;**

serve

309 classifier = medium.

310

311 - ****Hsu, Yu & Cheng ? trajectory + swing fusion**** (Sensors 2024, 24(13):4372,

312 [10.3390/s24134372](https://www.mdpi.com/1424-8220/24/13/4372)). Fusing TrackNet

trajectory with a

313 YOLOv7 swing detector via a Shot Refinement Algorithm raises hit detection ****F1 72.3%**

? 90.5%******

314 (over 69 BWF matches / 1582 rallies). Hits are found by ****trajectory direction-**

reversal / peak

315 detection****** (convert (x,y)?(frame,y) wave peaks) then reconciled with the swing cue.

****Why it fits:********

316 validates our ****late/score-level fusion**** approach in-domain, and direction-

reversal/peak is a cheap

317 ****per-window feature**** marking contact events (a "rally has hit cadence" signal ? the

discriminator

318 a "shuttle-moving" heuristic lacks). ****? Important caveat (the one 2-1 vote):****** this**

paper does

319 **within-rally stroke/hit* segmentation, ****not**** rally start/end ? "it splits merged*

rallies" is **our*

320 extrapolation*; frame it as "hit cadence plausibly helps locate rally boundaries," not

a proven

321 rally-splitter. ****Effort/payoff:****** the peak/direction-reversal feature = low-effort**

win (feed the

322 scorer); full swing fusion = medium.

323

324 - ****What best separates "rally" from "shuttle merely moving"****** (synthesis): sustained-**

inactivity

325 termination (See), hit cadence / direction-reversals (Hsu), net-crossings ?2 (our

Gate-A,

326 [`rally_gate.py`](../backend/pipeline/detectors/rally_gate.py)), two-sided court

occupancy (our

327 person features). None alone is sufficient; the spine (\$4) says ****feed them all to the**

learned

328 scorer******.

329

330 **### 5.3 Small-data strategies ? most quality per label**

331 - ****Timestamp / sparse supervision**** (Li et al., CVPR 2021,

[arXiv:2103.06669](https://arxiv.org/abs/2103.06669);

332 TAS survey Ding et al., [arXiv:2210.10352](https://arxiv.org/pdf/2210.10352)). Label

****one frame per**

333 action****** ? iteratively refine pseudo frame-labels ? **comparable** to full supervision

(50Salads:

334 timestamp 75.6% acc / 60.1 F1@50 vs full 83.7% / 70.1). ****Why it fits:****** sparse per-**

rally timestamps

335 cost far less than dense boundary labels ? relevant if we ever train a frame-wise model.

336 ****? Caveats:**** "comparable" hides a real ****~8?10 pt residual gap****; annotation savings are modest

337 (~35%, annotators still watch the whole clip); benchmarks are cooking/egocentric ? transfer to 1-D

338 shuttle trajectories is ****unproven****. ****Effort/payoff:**** research bet to pilot.

339 - ****Parameter-light scorer**** ? see TemporalMaxer (§5.1): keep the numpy LogisticRegression; resist a

340 learned temporal net until the corpus is much larger. ****Low-effort win (a discipline).****

341 - ****? OPEN (the research could NOT verify these ? honest gap):**** ****active learning*** (which video/segment

342 to label next via uncertainty/diversity/core-set) and the ****teacher-student / distillation governance**

343 boundary* for using Gemini as a weak/pseudo-labeler returned ****no surviving verified claim****. Our

344 pragmatic, guardrail-safe stance (not from the literature ? from our own design): pick the next

345 video to label by ****maximal local-vs-Gemini disagreement**** (the §4.4 shadow signal) + ****scorer**

346 uncertainty** + ****venue/fps diversity****; and treat ****Gemini strictly as an inference-time**

347 weak-reference / eval oracle ? never a training-data source for owned weights** (§8). These two

348 topics deserve their own grounded research pass before we lean on them.

349

350 **### 5.4 Evaluation rigor at small n ? lead with segmental metrics**

351 - ****Frame accuracy (MoF) and plain tIoU are BLIND to over-segmentation**** (TAS survey, Ding et al.,

352 [arXiv:2210.10352](https://arxiv.org/pdf/2210.10352); segmental-F1 origin Lea et al., CVPR 2017).

353 "The score can be high even when segments are fragmented." The over-merge that bit us is exactly

354 what tIoU-F1 hides (a giant window still ****overlaps* real rallies**). ****Action:**** lead the leaderboard

355 with the ****segmental Edit (Levenshtein) score**** (penalizes spurious ordering/fragmentation) ****and**

356 **F1@{0.1,0.25,0.5}****, paired with our existing `segment_count_ratio` + the merge/split breakdown.

357 ****We already have**** F1@k, mAP@tIoU, segment-count-ratio, bootstrap CIs, paired sign-test, Platt/

358 isotonic + ECE/Brier ? the only ****missing* primary metric** is the ****Edit score**** (low-effort add).

359 - ****? OPEN (not covered by a verified claim):**** IAA / label-noise level in the golden set, and which

360 over-segmentation-robust loss (ASRF boundary/smoothing, MS-TCN truncated-MSE) best tolerates it at

361 small n. Measure IAA as the corpus grows (Roora has an A/A overlap plan); revisit robust losses

362 only when a frame-wise model is on the table.

363

364 **### 5.5 Refuted claims (recorded so we do NOT repeat them)**

365 1. ****ASRF's boundary branch is model-independent ? a post-hoc splitter applicable to any dense**

366 segmentation output.**

367 ****REFUTED (0-3).**** Implication: we cannot bolt ASRF's boundary head onto

368 our trajectory signal as a drop-in; the two-branch model must be trained together. ? port the

369 ****concept*, don't expect a plug-in.**

369 2. ****A GFL-inspired Boundary Distribution Regression head is the mechanism most relevant to splitting**

370 merged windows.**

370 ****REFUTED (0-3)**** (and the cited source had a suspect identifier). Do not adopt

371 this framing as established.

372

373 ### 5.6 External due-diligence review (2026-06-17) ? owner-commissioned; verify before submission

374 An owner-commissioned external review (31 cited sources) of the

375 RALLY_DETECTION_QUALITY_REPORT.md checkpoint

376 **validated the methodology** (the decoupled 2-layer design, the R2 metric, LOVO + paired-sign-test discipline, the

377 consented-data + over-seg-aware-harness moat) and surfaced these **borrowable inputs**. These are

378 *research inputs for later pickup* ? they add **no measured numbers**, and the citations are the

379 review's own and **must be re-verified before any submission**.

380

381 - **Positioning ? the "first-mile" gap (highest-value).** Badminton stroke-classification SOTA

382 (**BST** ? Badminton Stroke-type Transformer, CVPR-W 2026; **DSTA** ? Decoupling Spatio-Temporal

383 Adapter) and the big datasets (**ShuttleSet** ~36,492 strokes; **Fine-Badminton** ~27,597 actions)

384 all **silently assume a pre-trimmed rally clip already exists** ? they depend on a human having

385 segmented the untrimmed match. Our system solves that ignored *prerequisite*. Reframe the paper as

386 ***"the missing untrimmed-video preprocessing pipeline that lets BST/MonoTrack run on raw footage"*** ?

387 a real, unoccupied gap and a far stronger narrative than "applied/systems."

388 - **R2 ? Action-Spotting / Precise-Event-Spotting (AS/PES).** Our asymmetric tolerance-match metric

389 sits at the **TAL × AS intersection** (interval task, tolerance-based match). Grounding it in the

390 AS/PES literature (SoccerNet, TennisSet, T-DEED, the PES survey [arXiv:2505.03991]) frames R2 as a

391 deliberate methodological contribution, not an ad-hoc metric ? strengthens §5.4 and the report's §9.

392 - **Comparative baselines reviewers will require.** Benchmark against **MonoTrack** and **BST** on

393 **ShuttleSet / Fine-Badminton** ? we currently have only a generic heuristic + the Gemini wrapper;

394 domain-specific baselines are a hard publication requirement.

395 - **Scorer upgrade ? prefer TemporalMaxer.** For the post-logreg scorer (corpus-gated; NEXT_STEPS M0.4),

396 **TemporalMaxer** (parameter-free temporal max-pool; ~2.8× fewer GMACs; robust at small-N,

397 [arXiv:2303.09055]) fits the cheap/local/fast mandate better than attention-heavy ActionFormer.

398 - **Court polygons ? pose ? the rally-END lever.** Court masking is not only the person-cue spectator

399 fix (§4 / MULTI_SIGNAL_FUSION_PLAN.md): once pose is isolated to on-court players it can resolve

400 **ambiguous, occluded end-of-rally states** ? feeding directly into the **#1 remaining gap (the

401 rally-END boundary)**, not just precision. Connect court-polygons to the R4/END-trim work.

402 - **Name the eval?serve lesson "pipeline distribution skew."*** Downstream-cue value depends on the

403 upstream proposal distribution (the Gate-A episode) ? a citable framing and candidate paper "lesson."

404 - **Citable Layer-1 anchor (public benchmark, NOT our footage).** WASB SBDT F1 95.6 (step=1) / 94.0

405 (step=3) vs MonoTrack 92.1, TrackNetV2 89.4, BallSeg 71.7 (WASB, BMVC 2023). Justifies the frozen

406 detector choice; **keep distinct from our own, still-unmeasured footage domain-gap** (§2.1, a real


```

407     open limitation the review also flags).
408
409     **Sequencing:** these are *publication-prep + post-current-idea* levers. The
prerequisites remain
410     corpus growth to n?16 and **multi-court isolated as its own reported slice** (§6 Tier-3
/ §7) ? do
411     those first; the items above land once current windowing/fusion ideas are exhausted.
412
413     ---
414
415     ## 6. Prioritized roadmap (cheap ? durable)
416
417     Every item is **its own measured PR**: default-OFF, scored on the §4 metric set over the
golden set
418     with honest stats (CI + sign-test), **re-confirmed on the served windowing** (the
eval?serve gate),
419     logged in [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md). Ordered by the research's
420     effort/confidence tiering. Each carries a **falsifiable success criterion** ? if it
isn't met, the
421     item does not promote (and that null is itself recorded).
422
423     ### Definition of "bridging the gap"
424     Local rally segmentation, measured **honestly LOVO-by-match on the golden set**,
approaches the
425     Gemini reference (today **0.93 single-court / 0.66 multi-court** tIoU-F1) ? but tracked
primarily
426     through the **over-segmentation-sensitive metrics** (segmental Edit + F1@k + low
merge/split rate),
427     since those, not bare tIoU-F1, capture our actual failure. A secondary, free signal is
rising
428     **agreement-with-Gemini** on the unlabeled pool (§4.4, shadow only).
429
430     ### Tier 0 ? Make the failure measurable FIRST (do before any fix)
431     - **E1 · Over-segmentation-aware leaderboard.** Add the **segmental Edit (Levenshtein)
score** to
432     `segmentation_metrics.py`; surface Edit + F1@{0.1,0.25,0.5} + `segment_count_ratio` +
a
433     **merge-rate / split-rate** breakdown (promote the
`scratch/compare_local_vs_gemini.py` matching
434     into a reusable `backend/eval/` helper) in the default `experiment.compare()` report;
add the
435     one-command rally-segmentation eval entrypoint (§4.5). **Success:** the current
heuristic's
436     over-merge shows up as a *number* (low Edit, high merge-rate) on the 6 golden videos,
and a
437     baseline row lands in the leaderboard. *(Low effort, high confidence ? research-backed
§5.4.)*
438
439     ### Tier 1 ? Low-effort, high-confidence wins
440     - **W1 · Bounded windowing (the direct over-merge fix).** Add a `max_window_duration`
cap and a
441     **sustained-inactivity split** to `trajectory_to_action_windows` ? port See et al.'s
rule (split /
442     end a window when ~80% of consecutive frames over a trailing window show no shuttle
motion).
443     Config-gated, default preserves current behavior. **Success:** segmental Edit + F1@k
improve vs
444     `CONTROL` on the golden set (?F1 CI excludes 0 **and** sign-test agrees) with **no
served-windowing
445     regression**; merge-rate drops materially. *(Win ? §5.2 See et al., ~95% end-of-rally
acc.)*
446     - **W2 · Rally-state features into the learned scorer (no special-casing).** Wire the
already-built,
447     default-OFF cues as scorer feature columns: **net-crossings** (Gate-A), **two-sided
motion**, and a

```

```

448     **new cheap trajectory feature ? direction-reversal / peak count** (Hsu) as a hit-
cadence proxy;
449     audio hit-rate when the hit detector lands. **Re-measure with W1 in place and on the
served
450     windowing** ? recall person was ?0.021 and audio +0.079 but *both CIs spanned 0* at
n=6 on the
451     *pre-fix over-merged* candidates; the picture may change once windows are bounded.
**Success:** the
452     ablation (`ablation.py`) shows a cue with ?F1 CI clearing 0 + sign-test, surviving
served re-check
453     ? promote that cue; others stay default-OFF + retained (no deletion on a small-n
null). *(Win/medium.)*
454     - **W3 · Keep the scorer parameter-light (a discipline, documented not coded).** Per
TemporalMaxer's
455     ablation (learned temporal conv overfits; parameter-free wins), resist replacing the
numpy
456     LogisticRegression with a learned temporal net until the corpus is much larger. Record
this as a
457     standing decision. *(Win ? $5.1 TemporalMaxer.)*
458
459     ### Tier 2 ? Medium effort
460     - **M1 · Serve-start classifier ? rally-start boundary** (See et al. pose/court, 88.6%).
Needs the
461     person cue + a **court-mask source** (manual polygon vs auto-homography ? *owner-
gated* per
462     [`MULTI_SIGNAL_FUSION_PLAN.md`](MULTI_SIGNAL_FUSION_PLAN.md) §9.1). **Success:**
start-boundary
463     timing error drops on golden; promotes only if it lifts segmental F1@k without served
regression.
464     - **M2 · Per-venue threshold calibration** against golden labels (the heuristic's
footage-fragility ?
465     it scored 0.157 on multi-court). Tune `velocity_threshold` + `merge_gap` so local
segment counts
466     track golden counts per venue. **Success:** multi-court F1 rises toward the Gen-0
head's 0.561
467     without harming clean-court F1. *(Velocity is already fps-normalized ? calibrate per-
venue, not
468     per-fps.)*
469
470     ### Tier 3 ? Research bets (durable; corpus-gated)
471     - **R1 · Owned segmenter with an explicit boundary branch** (the ASRF/TriDet *idea* ? a
boundary head
472     that predicts cut points ? ported into our owned TAL head, NOT the full deep model,
and **trained**,
473     not grafted (§5.5 refutation)). Per
[`OWNED_MODEL_IMPLEMENTATION_PLAN.md`](OWNED_MODEL_IMPLEMENTATION_PLAN.md):
474     **TAL head first, VLM later.** **Gate:** the central unverified assumption is whether
boundary-head
475     gains hold at our scale on trajectory features ? so **pilot only when the corpus is
materially
476     larger (~?16?30 matches)**, and only after W1/W2 establish the heuristic ceiling.
**Success:**
477     beats the tuned heuristic + LogReg on held-out matches (segmental Edit + F1@k),
honestly LOVO.
478     - **R2 · Timestamp/sparse supervision** to get more labeled rallies per annotator-hour
($5.3) ? pilot
479     if/when we train a frame-wise model; research bet (residual gap, unproven on
trajectories).
480
481     ### Throughout ? corpus growth + active learning + label quality
482     - **Grow the golden set 6 ? ~16** via Roora
([`GOLDEN_SET_VENDORING.md`](GOLDEN_SET_VENDORING.md)) ?
483     the single highest-leverage lever ([`NEXT_STEPS.md`](NEXT_STEPS.md)).
484     - **Active learning:** select the next videos to label by max **local-vs-Gemini
disagreement** (§4.4)

```

```

485 + scorer uncertainty + venue/fps diversity (our pragmatic stance; §5.3 flags the
formal methods as
486 an open research pass).
487 - **Measure IAA / label noise** as the corpus grows (Roora A/A overlap); revisit over-
segmentation-
488 robust losses only when a frame-wise model is on the table (§5.4 open question).
489
490 ### Dead-ends / do-NOT (so we don't relitigate)
491 - Don't treat **causal attention** (CausalTAD) as the over-merge fix ? it has no
boundary-split (§5.1).
492 - Don't expect **ASRF's boundary branch as a post-hoc drop-in** on our dense signal ?
refuted (§5.5).
493 - Don't adopt a **learned temporal net** at n=6?16 ? TemporalMaxer evidence says it
overfits (§5.1).
494 - Don't chase **bare tIoU-F1 / frame accuracy** ? blind to the over-merge (§5.4).
495 - Don't let **Gemini outputs become training data** for owned weights ? inference/shadow
only (§8).
496
497 ---
498
499 ## 7. Project: plan, tracking & definition of done
500
501 This roadmap is run as **one tracked project ? "Bridge the local?Gemini rally-detection
gap."** It
502 **completes when all committed tiers land AND a single quality number is attributed to
the effort**,
503 measured honestly and logged. This §7 is the living burn-down (update the status column
as PRs merge;
504 log each result in [`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md), numbers in
505 `eval_baselines/experiment_leaderboard.json`).
506
507 ### 7.1 The number (what we attribute to this effort)
508 Measured **leave-one-video-out, grouped by match, on the golden set** (grown to ?~16):
509 - **Headline:** multi-court rally-segmentation **tIoU-F1**, `baseline ? final`, plus
**gap-closure**
510 = `(final ? baseline) / (Gemini ? baseline)`. Today: heuristic **0.157** ? Gemini
**0.66**
511 (Gen-0 head already **0.561**). **Success bar (proposed, owner-adjustable): close ?
50% of the
512 heuristic?Gemini multi-court gap** ? i.e. local multi-court F1 ? ~0.41, ideally
beating Gen-0's 0.561.
513 - **Primary lens (over-segmentation-sensitive, the actual failure):** segmental
**F1@0.25** +
514 **Edit-score** + **merge-rate** ? these must improve even if tIoU-F1 moves little.
515 - **Reported with honest stats** (bootstrap CI + paired sign-test); a result counts only
if it also
516 holds on the **served** windowing (eval?serve gate). The **baseline is frozen at E1**
before any fix.
517 - **Tracked headline = the n=6 golden *aggregate* tIoU-F1@0.5** (LOVO; the robust unit,
not a single
518 video) via `python -m backend.eval.rally_seg_eval`. The 0.157/0.66 figures above are
per-venue
519 *references*; computing an exact "% gap-to-Gemini closed" needs Gemini measured on the
*same* golden
520 set + metric ? **TODO** (don't mix the per-venue refs with the aggregate).
521
522 **Progress (cumulative, golden n=6 aggregate tIoU-F1@0.5):**
523
524 | after | F1 | merge_rate | ? vs baseline |
525 |---|---|---|---|
526 | E1 baseline (current heuristic) | 0.300 | 0.523 | ? |
527 | + W1 bounded windowing (cap=12, **default-OFF**) | **0.439** | **0.038** | **+0.139**
(6/0 sign-test) |
528 | + W2 net-crossings gate (mc=1, **default-OFF**) | **0.457** | 0.038 | **+0.157
cumulative** (W2 step +0.019, 6/0) |

```

529 | + W1.5 hard-cap fallback (**default-OFF**) | 0.444 | 0.000 | **+0.006 (n.s., CI spans 0)** ? neutral on golden; real-footage fix (mahadevpura-1 174s blob ? 24x?9s windows) |

530 | + W2 multi-feature scorer (person/audio) | _next (GPU)_ | | |

531

532 ### 7.2 Work items (status checklist)

533 Status: ? not started · ? in progress · ? done. (Each row = its own default-OFF, measured PR.)

534

535 | ID | Deliverable | Tier | Depends on | Acceptance / promote criterion | Status |

536 |---|---|---|---|---|---|

537 | **E1** | merge/split-rate metric + F1@k surfaced + one-command `rally\_seg\_eval` entrypoint; **froze baseline** (#185) | 0 | ? | over-merge shows up as a number on the 6 golden videos | ? baseline F1 0.300, merge_rate 0.523 |

538 | **W1** | Bounded windowing ? `max\_window\_duration` cap + inactivity split (See et al.) (#187) | 1 | E1 | seg F1@k ? vs baseline (CI?0 **and** sign-test), no served regression; merge-rate ? | ? **?F1 +0.139 (6/0), merge_rate 0.523?0.038** (default-OFF) |

539 | **W1.5** | Hard-cap fallback (`window\_hard\_cap`) ? chop continuously-active windows at the cap when no inactivity gap exists | 1 | W1 | golden ?F1 ?0 (or neutral) **** fixes a real-footage over-merge it can't otherwise cap | ? **golden ?F1 +0.006 (n.s.)**; real `mahadevpura-1` 174s blob ? 24x?9s; retained **default-OFF** (real-footage + true-cap rationale, not a golden lift) |

540 | **W2** | Rally-state features (net-crossings ? #188 + **never-empty safeguard** ?; two-sided motion, **direction-reversal/peak**, audio = follow-on) | 1 | E1, W1 | ?1 cue ?F1 CI clears 0 + sign-test, survives served re-check ? promote; others retained default-OFF | ? **net-crossings ?F1 +0.019 (6/0)**; gate now **do-no-harm** (never zeroes a video); multi-feature scorer next (GPU) |

541 | **W3** | Keep scorer parameter-light (standing decision; TemporalMaxer evidence) | 1 | ? | documented; no learned temporal net until corpus ? now | ? |

542 | **M1** | Serve-start classifier ? rally-start boundary (pose/court) | 2 | W1, **court-mask decision (owner-gated)** | start-boundary timing error ?; promote on seg-F1 lift, no served regression | ? |

543 | **M2** | Per-venue threshold calibration vs golden | 2 | E1 | multi-court F1 ? toward Gen-0 0.561 without harming clean-court | ? |

544 | **R1** | Owned segmenter with an explicit **boundary head** (port ASRF/TriDet *idea*, trained ? not grafted) | 3 (stretch) | corpus ?~16?30, W1/W2 | beats tuned heuristic+LogReg on held-out matches (seg Edit + F1@k), LOVO | ? |

545 | **R2** | Timestamp/sparse-supervision pilot (label-efficiency) | 3 (stretch) | R1 path | measured label-efficiency win | ? |

546 | **C** | Golden corpus 6 ? ~16 (Roora) + active-learning selection + measure IAA | ongoing | ? | corpus ?16; IAA recorded | ? (~6 now; ~10 by the weekend) |

547

548 ### 7.3 Sequencing

549 `E1` (metrics + baseline) ? `W1`/`W2`/`W3` (the high-confidence wins) ? `M1`/`M2` ? `R1`/`R2`

550 (corpus-gated). `C` (corpus growth) runs in parallel throughout ? it is the single highest-leverage

551 input and gates the Tier-3 bets.

552

553 ### 7.4 Definition of done

554 The project is **complete** when:

555 1. **Tier 0 + Tier 1 + Tier 2** work items are merged ? each a measured PR (default-OFF, promoted only

556 on measured lift, re-confirmed on the served windowing).

557 2. **The headline number (\$7.1) is recorded** in `QUALITY\_ITERATIONS.md` + the leaderboard, with honest

558 stats, on the grown golden set ? **and the gap-closure target is met, or the honest null is recorded

559 with a reasoned go/no-go on Tier 3.**

560 3. Each landed item left the harness greener (no regression) and `QUALITY\_ITERATIONS.md` tells the story.

561

562 > **Tier-3 scope decision (owner-adjustable):** R1/R2 (the owned boundary-head model + sparse

563 > supervision) are **corpus-gated stretch goals** ? the verified research says deep TAL detectors risk

```

564     > overfitting at n=6?16. **Default:** the project *completes* at Tier 0?2 + an
attributed number, and
565     > Tier 3 spins out into the owned-model milestone
(['OWNED_MODEL_IMPLEMENTATION_PLAN.md'](OWNED_MODEL_IMPLEMENTATION_PLAN.md)).
566     > Fold Tier 3 into the DoD instead if you want the project to span the owned model.
567
568     ---
569
570     ## 8. Constraints & guardrails (non-negotiable) + cross-links
571
572     - **Licensing:** every dependency ? code, data, AND weights ? must be **MIT / Apache-2.0
/ BSD**.
573         Reject viral/NC/custom (Gemma-3, Llama-4 MAU cap, Commons-Clause).
574     - **Never train owned weights on paid-LLM (Gemini/OpenAI/Anthropic) outputs**
(terms/copyright);
575         paid LLMs are inference/serving/shadow-reference only. Also forbidden: public datasets
known to
576         be LLM-generated. **Exclude minors** from any training pool; per-user data needs
explicit opt-in.
577     - **Owned model:** TAL head FIRST (de-risked, trains in minutes, doubles as a pseudo-
label teacher +
578         honest baseline the VLM must beat), then Qwen3-VL (Apache-2.0) / InternVL3 (MIT) ? see
579         ['OWNED_MODEL_IMPLEMENTATION_PLAN.md'](OWNED_MODEL_IMPLEMENTATION_PLAN.md).
580     - **Cross-links:** ['HOW_RALLY_DETECTION_WORKS.md'](HOW_RALLY_DETECTION_WORKS.md) .
581         ['EXPERIMENT_HARNESS.md'](EXPERIMENT_HARNESS.md) .
['MULTI_SIGNAL_FUSION_PLAN.md'](MULTI_SIGNAL_FUSION_PLAN.md) .
582         ['ANALYZERS_AND_RECONCILERS.md'](ANALYZERS_AND_RECONCILERS.md) .
['ABLATION_LAYER.md'](ABLATION_LAYER.md) .
583         ['EVALUATION.md'](EVALUATION.md) . ['QUALITY_ITERATIONS.md'](QUALITY_ITERATIONS.md) .
584         ['GOLDEN_SET_VENDORING.md'](GOLDEN_SET_VENDORING.md) .
['OWNED_MODEL_IMPLEMENTATION_PLAN.md'](OWNED_MODEL_IMPLEMENTATION_PLAN.md).
585

```

6. user (2026-06-20T08:00:36.515Z)

```

1     # Request-Access form ? segmented lead capture + setup discovery
2
3     > **Status:** `IN PROGRESS` . **Owner:** Avi . **Created:** 2026-06-20 . **Last
updated:** 2026-06-20
4     > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived`
5     > **Relation:** first tracked work item in the `khelsutra` repo; mirrors the engine
repo's
6     > `PROJECT_DOC_TEMPLATE` discipline (status header . decisions-locked . $6 tracker .
DoD).
7     > Feeds the platform direction in the engine repo's `PLATFORM_ARCHITECTURE.md` $3`
(StorageBackend /
8     > ComputeBackend) and `VIDEO_LOCALITY_MODEL.md` ? see $7.
9
10    ---
11
12    ## 0. TL;DR
13    Replace the landing page's `mailto:` "Request access" CTA with a real form that (a)
segments
14    **academies vs individual enthusiasts**, (b) asks the two questions that actually decide
how we'd
15    serve a user ? **"do you have a GPU machine?" (compute placement) and **"where do your
videos
16    live?" (storage backend) ? and (c) keeps a free-text message. Submissions are stored
in our own
17    **Cloudflare D1** via a small **Worker** (`POST /api/request-access`). The form doubles
as PMF +
18    architecture research.
19
20    ## 1. Problem & goal
21    Today the CTA is `mailto:hello@khelsutra.guru` ? zero structure, nothing collected, no

```

```

segmentation.
22     We want **structured, queryable** interest signals from two distinct audiences, captured
the moment
23     someone raises a hand, **without shipping lead data to a third party** (privacy-first
ethos).
24     **Goal:** a clean on-brand form whose responses land in our own Cloudflare account,
queryable later.
25
26     ## 2. Decisions locked
27     - **Backend = Cloudflare Worker + D1** (owner decision, 2026-06-20). Data stays in our
Cloudflare
28     account; queryable; no third party. (Alternatives considered: hosted form
(Tally/Formsfree ?
29     third-party data), Worker?email-only (no DB), structured `mailto` (no collection).
Rejected for
30     the reasons in parens.)
31     - **Two audience segments:** `academy` and `individual` ? the form lightly adapts
(academy also
32     asks club name + #courts).
33     - **The two discovery questions are first-class** (not optional polish): GPU/compute +
storage ?
34     they map 1:1 onto the BY0-compute / BY0-storage architecture forks ($7).
35     - **Free-text message retained** ("Anything else?").
36     - **Privacy:** store only what's submitted + a coarse country (from
`request.cf.country`); **no raw
37     IP**; consent checkbox for "email me about the alpha"; this data is
**operations/contact only**,
38     never a training source (same guardrail as the engine repo).
39     - **Deploy model changes:** the site moves from dashboard direct-upload ? **Wrangler**
(Worker with
40     static assets + D1 binding). Documented in $5; owner runs the `wrangler` deploy (our
account/auth).
41
42     ## 3. Form spec
43     | Field | Type | Notes |
44     |---|---|---|
45     | Segment | radio | `academy` \| `individual` (required) |
46     | Name | text | required |
47     | Email | email | required (so we can reply) |
48     | City | text | optional |
49     | Club / academy name | text | shown for `academy` |
50     | Courts recorded | select | `1` \| `2-4` \| `5+` (academy; optional) |
51     | **GPU machine for local processing?** | radio | `yes` \| `unsure` \| `no` \| `prefer-
cloud` |
52     | **Where do your recordings live / how would you share them?** | checkboxes (multi) |
`gdrive` \| `onedrive` \| `dropbox` \| `bucket` (S3/GCS) \| `local` (SD/drive) \| `unsure` |
53     | Message | textarea | free text ("Anything else?") |
54     | Consent | checkbox | "Okay to email me about the KhelSutra alpha" |
55     | `hp` | hidden honeypot | anti-spam; if filled ? silently drop |
56
57     ## 4. Design / architecture
58     - **Frontend** (`site/index.html`): the header + hero "Request access" buttons scroll to
the
59     `#join` form; `#join` becomes the form. On submit, JS `POST`s JSON to `/api/request-
access`.
60     **Progressive enhancement:** if the endpoint isn't live yet (pre-Worker), it falls
back to a
61     prefilled `mailto:hello@` so the form is useful from the first re-upload.
62     - **Worker** (`/api/request-access`): validates (email present, honeypot empty), inserts
into D1,
63     returns `200 {ok:true}`; serves the static assets for all other routes.
64     - **D1 schema** (`access_requests`): `id, created_at, segment, name, email, city,
org_name, courts,
65     has_gpu, storage (csv), message, consent, country, user_agent`.
66     - **Anti-spam:** honeypot first; add **Cloudflare Turnstile** in P2 if spam appears.

```

```

67     - **Notify (P2):** optional email ping to `hello@` per submission (MailChannels /
Cloudflare Email).
68
69     ## 5. Deploy (owner-run; our Cloudflare account)
70     ```
71     wrangler d1 create khelsutra-access          # once; copy the database_id into
wrangler.toml
72     wrangler d1 execute khelsutra-access --file=site/schema.sql # create the table
73     wrangler deploy                             # ships the Worker + static assets + D1
binding
74     ```
75     > Until the Worker ships, the form works via the `mailto` fallback (re-upload `site/` as
today).
76     > Git?Cloudflare auto-deploy still intentionally not connected.
77
78     ## 6. Deliverables & progress tracker ? source of truth
79     Legend: ? Todo · ? In progress · ? Done. One PR per row off `origin/master`.
80     | ID | Deliverable | Gated? | Status | PR |
81     |----|-----|-----|-----|----|
82     | P0 | This design doc + `NEXT_STEPS.md` | No | ? | #2 |
83     | P1 | Form UI (segments + GPU + storage + free text + consent) w/ `mailto` fallback |
No | ? | #2 |
84     | P2 | Worker + D1 + `wrangler.toml` + `schema.sql` + **unit tests (10 green)** | owner
deploy | ? | this PR (built+tested; owner runs `wrangler d1 create` + `deploy`) |
85     | P3 | Anti-spam (Turnstile) + per-submission email notify | No | ? | ? |
86     | P4 | Tiny admin/export view (auth'd) to read responses | owner | ? | ? |
87
88     ## 7. Why this matters beyond lead capture
89     The two questions are **the architecture forks** from the engine repo's BYO analysis:
90     - **GPU?** ? compute placement: desktop/local (their GPU) vs hosted vs zero-infra-BYO
91       (`PLATFORM_ARCHITECTURE.md §3.2`, `DESKTOP_APP_PLAN.md` P0 spike).
92     - **Storage?** ? which `StorageBackend` + whether autosync is worth building
93       (`PLATFORM_ARCHITECTURE.md §3.1`, the cloud-storage/autosync thread).
94     So early responses directly inform what to build next ? not just who to email.
95
96     ## 8. Open questions
97     1. Turnstile now or wait for spam? (Default: honeypot now, Turnstile if needed.)
98     2. Email-notify per submission, or just query D1 weekly? (Default: notify in P2.)
99     3. Do we localize the form (Kannada/Hindi) for academies now or at beta? (Default:
English now.)
100
101     ## 9. Definition of done
102     - [ ] Form replaces the `mailto` CTA, segments academy/individual, asks GPU + storage +
free text.
103     - [ ] Submissions persist to Cloudflare D1 via the Worker; honeypot drops bots.
104     - [ ] Privacy posture honored (no raw IP; consent; contact-only).
105     - [ ] Deployed live on khelsutra.guru; a test submission appears in D1.
106     - [ ] Doc tracker + DoD updated; archived when P1?P3 land.
107

```

7. user (2026-06-20T08:00:36.552Z)

```

1     # Portability ? Cloudflare now, always-on VM if needed
2
3     > **Status:** design + proven (2026-06-20). **Decision:** keep the site/form host-
portable so
4     > moving off Cloudflare to an owner-run always-on VM ("if Cloudflare acts out") is an
**adapter
5     > swap, not a rewrite.**
6
7     ## The shape that makes it portable
8     - **Logic is runtime-agnostic.** `src/index.js` uses Web-standard `Request`/`Response`
and pure
9     helpers (`validate`, `toRow`, `isBot`, `handleRequest`) ? no host-specific calls.
10    - **All DB coupling is in adapters** (`src/store.js`): `d1Store` (Cloudflare D1) and

```

```

`sqliteStore`
11     (node:sqlite / better-sqlite3). `handleRequest(request, store)` takes the store by
injection.
12     - **One schema, one SQL.** `schema.sql` and the `INSERT` are plain SQLite, so D1 and a
VM's SQLite
13     file are byte-for-byte compatible ? zero dialect change. (Postgres later = minor
placeholder tweaks.)
14
15     ## Two hosts, same code
16     | | Cloudflare (now) | Always-on VM (fallback) |
17     |---|---|---|
18     | Entry | `src/index.js` (Worker `fetch`) | `server.js` (node:http) |
19     | Static | `env.ASSETS` | `server.js` serves `public/` |
20     | Storage | `d1Store(env.DB)` ? D1 | `sqliteStore(db)` ? `node:sqlite` file |
21     | Country | `cf-ipcountry` header | optional (GeoIP / null) |
22     | Deploy | git ? Workers Builds | `node server.js` under systemd |
23     | TLS/DNS | Cloudflare | keep Cloudflare in front, **or** Caddy/Let's Encrypt |
24
25     **Enforced, not asserted.** `node --test src/` runs BOTH paths and CI
(`.github/workflows/test.yml`,
26     Node 24) runs them on every PR, so the repo stays compatible with both hosts:
27     - `src/index.test.js` ? the Cloudflare Worker + D1 logic.
28     - `src/server.test.js` ? **boots the real `server.js`** (node:http + node:sqlite) and
round-trips it
29     over HTTP into SQLite (the same proof we run on the droplet, automated).
30
31     For a live check against a real VM: `SSH_TARGET=root@host scripts/droplet-smoke.sh` (scp
+ boot +
32     round-trip + cleanup). Verified live on the owner's droplet (Ubuntu 24.04, isolated
`/srv/khelsutra-site`).
33
34     ## VM deploy runbook ("if Cloudflare acts out")
35     Node 22+ only ? **no `npm install`** (node:sqlite is built in).
36     ```
37     # on the VM:
38     git clone ? && cd khelsutra/site          # or scp the site/ folder
39     PORT=8080 DB_PATH=/var/lib/khelsutra/access.db node server.js
40     # put it under systemd (auto-restart) + a reverse proxy (Caddy/nginx) for TLS,
41     # then point DNS at the VM. ~1 hour; logic + schema unchanged.
42     ```
43
44     ## Honest notes
45     - "Cloudflare acts out" usually means **Workers/D1 pricing or limits, not DNS** ? so the
cheapest
46     move is to keep Cloudflare DNS/CDN in front and shift only the **compute** to the VM
(point a
47     route/origin at the box). Full exit (own DNS + Caddy TLS) is also supported.
48     - You can run **both** (Cloudflare primary + VM warm standby) since the schema is
identical; a
49     periodic `access_requests` export keeps them in sync.
50     - `server.js` is intentionally dependency-free; if you later want Postgres or a
framework (Hono runs
51     on Workers **and** Node), it slots in behind the same `store` + handler.
52

```

8. assistant (2026-06-20T08:00:39.647Z)

I have everything I need from the canonical template and the three real-world example docs. Let me compile the exact schema.

9. user (2026-06-20T08:01:27.691Z)

Structured output provided successfully

10. assistant (2026-06-20T08:01:32.191Z)

I've extracted the exact tracked-design-doc structure and returned it via StructuredOutput. The schema is grounded in the canonical template (`docs/PROJECT\_DOC\_TEMPLATE.md`) and verified against three real in-use examples (`RALLY\_QUALITY\_RESEARCH.md`, `PROMINENT\_COURT\_DETECTION.md`, and the khelsutra `REQUEST\_ACCESS\_FORM.md` / `PORTABILITY.md`), with path:line citations for every field.