claude-windows / c47f0851-d75f-45b6-9126-e48544381426

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c47f0851-d75f-45b6-9126-e48544381426\subagents\workflows\wf_5af9a132-4ce\agent-a04b0e9965c4bdf90.jsonl`
- SHA-256: `0fa3fb1d64693b7fae65d0dbc9c285364e2ab38265c4449b00a3138ddf9c10e0`
- Source modified: `2026-06-20T09:02:23+00:00`
- Imported at: `2026-07-05T16:48:23+00:00`
- project: `wf_5af9a132-4ce`
- session_id: `c47f0851-d75f-45b6-9126-e48544381426`

Transcript

1. user (2026-06-20T09:00:58.224Z)

Adversarially verify this review finding against the ACTUAL code in the working tree. Default to 'false-positive' unless you can point to the concrete line(s) and a real failure scenario.
Context: Repo: badminton-highlight-indexer (Python). Branch feat/native-wasb-runner off origin/master.
The change builds the M3.5 Linux-native WASB detector runner + wires real extraction into the worker.
Files changed/added (read them in the WORKING TREE ? changes are uncommitted):
- backend/pipeline/detectors/native_wasb_runner.py (NEW ? the native runner)
- backend/pipeline/detectors/wasb_infer.py (added --device, threaded device into _build_cfg/run/run_video_streaming/main; device-keyed the resumable manifest)
- backend/pipeline/segmenters/trajectory_hybrid.py (_resolve_runner: new 'native' branch + _build_native_runner base hook; generalized 'WSL env' log lines)
- backend/pipeline/segmenters/wasb_hybrid.py (_build_native_runner override)
- backend/pipeline/segmenters/fusion_hybrid.py (_build_native_runner override; WASB substrate only)
- backend/config/models.py (WasbIndexConfig.device + .python_bin; detector_impl doc mentions 'native')
- backend/worker/__main__.py (cmd_train: --trajectory-source {synth,detector} + --segmenter; _resolve_train_detector; detector path pulls videos + runs runner)
- tests/test_native_wasb_runner.py, tests/test_worker.py (new tests)

DESIGN INTENT / CONTRACTS (judge against these):
- The EXISTING WSL runner is backend/pipeline/detectors/wasb_runner.py (WasbRunner, WslCommandMixin). The native runner must be a sibling that runs the SAME wasb_infer.py natively (no wsl, no 'conda activate' ? it invokes the wasb conda env python BY PATH, no /mnt path translation) and emit a BYTE-IDENTICAL Frame,Visibility,X,Y CSV with the SAME filename '{stem}__{vid12}_wasb.csv'.
- ACCEPTANCE GATE = byte-parity: same clip via WSL (Windows) vs native (Linux, device=cuda) -> identical CSV + identical candidate windows. device default 'cuda' must keep wasb_infer's Hydra overrides identical to before (parity).
- Selected via the detector_impl seam: indexing.detector_impl / RALLY_DETECTOR_IMPL = 'native'.
- Constraint: 'backend' must NOT import 'training' (worker shells out). The default trajectory source stays 'synth' so the CPU e2e is unchanged.
- The code RUNS on Linux (GPU box) but the TESTS run on the dev's Windows machine; cross-platform path handling matters.
- weights/repo/python/scratch come from env first (WASB_WEIGHTS/WASB_REPO_DIR/WASB_PYTHON/WASB_SCRATCH) then indexing.wasb.* then default.

FINDING [worker-extract] (major) @ backend/worker/__main__.py:232-234 (cmd_train detector branch ? video_by_stem build) and :244 (video lookup): video_by_stem is built from `storage.list_uris(f"{corpus}/videos/")` which (LocalStorage.list / object-store list) recurses and returns EVERY file under videos/, with NO extension filter ? unlike the labels side which filters to .csv (line 217). Any non-video sidecar (v1.srt, v1.json, v1.txt, a stray thumbnail)

is indexed by stem, and on a stem collision the dict silently keeps the LAST-listed entry.
WHY: Confirmed: for files [v1.json, v1.mp4, v1.srt] sorted (the order LocalStorage.list yields),
video_by_stem['v1'] ends up = 'v1.srt' ? the real v1.mp4 is shadowed. The runner then runs on a
subtitle/sidecar: compute_video_id hashes it, the WASB runner fails (or the stub blindly
'succeeds' on garbage bytes, polluting the training manifest with a trajectory for the wrong
file). Buckets routinely carry per-video sidecars, so this is a realistic silent corruption /
wrong-skip.
PROPOSED FIX: Filter the videos listing to known video extensions, mirroring the .csv filter:
e.g. `if storage.parse_uri(vuri).name.lower().endswith(('.mp4','.mov','.mkv','.avi'))` before
indexing; and on collision prefer a video extension / warn. At minimum restrict to the same
extensions as security.allowed_upload_extensions.

2. assistant (2026-06-20T09:01:00.061Z)

I'll verify this finding against the actual code in the working tree.

3. user (2026-06-20T09:01:01.215Z)

```
180             print(f"[worker] --segmenter {args.segmenter!r} has no shuttle detector; "
181                   f"use a trajectory hybrid (wasb_hybrid / tracknet_hybrid /
fusion_hybrid).")
182             return None
183         runner, _ = seg._resolve_runner(config.get("indexing", {}))
184         ok, msg = runner.healthcheck()
185         if not ok:
186             print(f"[worker] detector environment not ready: {msg}")
187             return None
188         print(f"[worker] detector ready ({args.segmenter}): {msg}")
189         return runner
190
191
192     def cmd_train(args: argparse.Namespace) -> int:
193         """Gen-0 train-from-bucket: pull labels (+ videos) -> trajectories -> manifest ->
shell out
194         to training.gen0.harness (backend must NOT import training) -> push the artifact
bundle.
195
196         Two trajectory sources (the train pipeline + harness are identical either way):
197           * ``synth`` (default) ? the CPU "mock GPU": deterministic, label-consistent
synthetic
198             shuttle paths (no GPU/WSL), so the whole pipeline is validated on a CPU box /
CI.
199           * ``detector`` ? REAL shuttle trajectories: pull each video and run the resolved
200             DetectorRunner (native WASB on a GPU box, or stub for a GPU-free wiring test).
This
201             is the M3.5 "first real training" path; only the trajectory source flips.
202         """
203         import subprocess
204
205         config = load_config(args.config) if args.config else load_config()
206         _apply_overrides(config, args.set)
207         storage_cfg = config.storage.model_dump()
208
209         scratch = args.scratch or tempfile.mkdtemp(prefix="rally-train-")
210         labels_dir = os.path.join(scratch, "labels")
211         traj_dir = os.path.join(scratch, "trajectories")
212         videos_dir = os.path.join(scratch, "videos")
213         os.makedirs(labels_dir, exist_ok=True)
214         os.makedirs(traj_dir, exist_ok=True)
215
216         corpus = args.corpus_uri.rstrip("/")
217         label_uris = sorted(u for u in storage.list_uris(f"{corpus}/labels/",
config=storage_cfg)
218                             if u.lower().endswith(".csv"))
219         if len(label_uris) < 2:
```

```
            print(f"[worker] need >=2 labeled videos for leave-one-video-out; found
{len(label_uris)} "
                  f"under {corpus}/labels/")
            return 2

        # Real-detector source: resolve the runner once + index the bucket's videos/ by
stem.
        runner = None
        video_by_stem: dict = {}
        if args.trajectory_source == "detector":
            os.makedirs(videos_dir, exist_ok=True)
            runner = _resolve_train_detector(args, config)
            if runner is None:
                return 2
            for vuri in storage.list_uris(f"{corpus}/videos/", config=storage_cfg):
                vname = storage.parse_uri(vuri).name
                video_by_stem[os.path.splitext(vname)[0]] = vuri

        print(f"[worker] trajectory source: {args.trajectory_source}")
        specs: List[dict] = []
        for uri in label_uris:
            fname = storage.parse_uri(uri).name                    # e.g.
GX020094_proxy.rallies.csv
            name = fname.replace(".rallies.csv", "").replace(".csv", "")
            local_label = storage.fetch(uri, os.path.join(labels_dir, fname),
config=storage_cfg)
            if args.trajectory_source == "detector":
                vuri = video_by_stem.get(name)
                if not vuri:
                    print(f"[worker] no video for {name!r} under {corpus}/videos/ ?
skipping.")
                    continue
                local_video = storage.fetch(vuri, os.path.join(videos_dir,
storage.parse_uri(vuri).name),
                                            config=storage_cfg)
                video_id = compute_video_id(local_video)
                # Real fps/width drive the trajectory frame->time mapping (synth fixes them
at 30/1280).
                from backend.pipeline.segmenters.trajectory_hybrid import
TrajectoryHybridSegmenter
                fps, frame_width = TrajectoryHybridSegmenter._probe_fps_width(local_video)
                traj = runner.run_predict(local_video, traj_dir, video_id=video_id)
                if not traj:
                    print(f"[worker] detector produced no trajectory for {name!r} ?
skipping.")
                    continue
                specs.append({"name": name, "trajectory": traj, "labels": local_label,
                              "fps": fps, "frame_width": frame_width})
            else:
                from backend.pipeline.detectors.stub_runner import
synth_trajectory_from_labels
                traj = os.path.join(traj_dir, f"{name}_traj.csv")
                synth_trajectory_from_labels(local_label, traj, fps=args.fps,
frame_width=args.frame_width)
                specs.append({"name": name, "trajectory": traj, "labels": local_label,
                              "fps": args.fps, "frame_width": args.frame_width})

        if len(specs) < 2:
            print(f"[worker] need >=2 videos with trajectories for leave-one-video-out; got
{len(specs)}.")
            return 2
        manifest_path = os.path.join(scratch, "manifest.json")
        with open(manifest_path, "w", encoding="utf-8") as f:
            json.dump(specs, f, indent=2)
        src_label = "real detector" if args.trajectory_source == "detector" else "CPU-mock
```

```
stub"
273         print(f"[worker] built manifest: {len(specs)} videos ({src_label} trajectories)")
274
275         out_dir = os.path.join(scratch, "artifacts")
276         cmd = [sys.executable, "-m", "training.gen0.harness",
277                "--manifest", manifest_path, "--out", out_dir, "--version", args.version]
278         if args.floor:
279             floor_local = (storage.fetch(args.floor, os.path.join(scratch, "floor.json"),
config=storage_cfg)
280                            if "://" in args.floor else args.floor)
281             cmd += ["--floor", floor_local]
282         print("[worker] training:", " ".join(cmd))
283         rc = subprocess.run(cmd).returncode
284         if rc != 0:
285             print(f"[worker] gen0 harness failed (rc={rc})")
286             return rc
287
288         # PUSH the versioned artifact bundle (weights.json + manifest.json) to the bucket.
289         bundle = os.path.join(out_dir, args.version)
290         out_prefix = args.out_uri.rstrip("/")
291         pushed = []
292         for fn in sorted(os.listdir(bundle)):
293             uri = f"{out_prefix}/weights/{args.version}/{fn}"
294             storage.put(os.path.join(bundle, fn), uri, config=storage_cfg)
295             pushed.append(uri)
296         print(f"[worker] train done: pushed {len(pushed)} artifact(s):")
297         for u in pushed:
298             print("   ", u)
299         return 0
```

4. assistant (2026-06-20T09:01:05.035Z)

I have the core code. Now let me verify the claims about `storage.list_uris`, `parse_uri`, and `LocalStorage.list` ordering behavior, plus the `.csv` filter and extension filters.

5. user (2026-06-20T09:01:06.000Z)

No files found

6. user (2026-06-20T09:01:06.317Z)

```
Found 8 files
backend\storage\s3.py
backend\storage\ref.py
backend\storage\gdrive.py
backend\storage\local.py
backend\storage\gcs.py
backend\storage\__init__.py
backend\storage\base.py
docs\PLATFORM_ARCHITECTURE.md
```

7. user (2026-06-20T09:01:09.626Z)

```
1       """Pluggable storage layer: local (default) + S3/GCS first-class buckets + Drive
(follow-up).
2
3       URI-driven (see ``ref.parse_uri``). The package import is cheap ? cloud SDKs are
imported
4       lazily, only when a non-local URI is actually used (install the ``storage-s3`` /
``storage-gcs``
5       / ``storage-gdrive`` extras). The default ``local`` backend needs no third-party deps
and
6       preserves today's behaviour byte-for-byte.
7
```

```
8       Worker/dispatcher entry points speak URIs via the thin helpers ``fetch`` / ``put`` /
9       ``list_uris`` / ``get_storage``; ``local://`` (and bare paths) make ``fetch`` an
identity
10      passthrough. See ``docs/DECOUPLED_COMPUTE/``.
11      """
12      from __future__ import annotations
13
14      import os
15      from typing import Any, Dict, Iterator, Optional
16
17      from backend.storage.base import StorageBackend
18      from backend.storage.ref import StorageRef, StorageStat, parse_uri
19
20      __all__ = [
21          "StorageBackend", "StorageRef", "StorageStat", "parse_uri",
22          "get_storage", "fetch", "put", "list_uris",
23      ]
24
25
26      def _backend_class(name: str):
27          """Lazy-resolve a backend class so cloud SDKs aren't imported until needed."""
28          if name == "local":
29              from backend.storage.local import LocalStorage
30              return LocalStorage
31          if name == "s3":
32              from backend.storage.s3 import S3Storage
33              return S3Storage
34          if name == "gcs":
35              from backend.storage.gcs import GCSStorage
36              return GCSStorage
37          if name == "gdrive":
38              from backend.storage.gdrive import GDriveStorage
39              return GDriveStorage
40          raise ValueError(f"unknown storage backend: {name!r} (expected
local|s3|gcs|gdrive)")
41
42
43      def get_storage(backend: Optional[str] = None, *,
44                      config: Optional[Dict[str, Any]] = None, **overrides) -> StorageBackend:
45          """Construct a backend. ``backend`` (or ``config['backend']``, default 'local')
selects
46          the implementation; its per-backend settings come from ``config[backend]`` merged
with
47          ``overrides``. Settings carry NO secrets ? only endpoints/regions/paths/source-
selectors."""
48          cfg = config or {}
49          name = backend or cfg.get("backend", "local")
50          settings: Dict[str, Any] = dict(cfg.get(name, {}) or {})
51          settings.update(overrides)
52          return _backend_class(name)(**settings)
53
54
55      def fetch(uri: str, dst: "str | os.PathLike[str] | None" = None, *,
56                config: Optional[Dict[str, Any]] = None, overwrite: bool = False) -> str:
57          """Resolve ``uri`` to a local file path. For ``local`` this is an identity
passthrough
58          (no copy); for s3/gcs it downloads to ``dst``."""
59          ref = parse_uri(uri)
60          return get_storage(ref.backend, config=config).fetch_to_local(ref, dst,
overwrite=overwrite)
61
62
63      def put(local_path: "str | os.PathLike[str]", uri: str, *,
64              config: Optional[Dict[str, Any]] = None, content_type: Optional[str] = None) ->
str:
```

```
65          """Upload/copy ``local_path`` to ``uri``; return ``uri``."""
66          ref = parse_uri(uri)
67          get_storage(ref.backend, config=config).put(local_path, ref,
content_type=content_type)
68          return uri
69
70
71      def list_uris(prefix_uri: str, *, config: Optional[Dict[str, Any]] = None) ->
Iterator[str]:
72          """Yield the URI of every object under ``prefix_uri``."""
73          ref = parse_uri(prefix_uri)
74          for r in get_storage(ref.backend, config=config).list(ref):
75              yield r.uri
76
```

8. user (2026-06-20T09:01:09.810Z)

```
1       """Local-filesystem storage backend ? the default, zero new dependencies.
2
3       Wraps today's behaviour byte-for-byte: ``fetch_to_local`` is an identity passthrough
(the
4       file is already local), ``put`` is a copy (a no-op when src == dst). With
``backend='local'``
5       the whole pipeline behaves exactly as before the storage layer existed.
6       """
7       from __future__ import annotations
8
9       import os
10      import shutil
11      from datetime import datetime
12      from typing import Iterator
13
14      from backend.storage.base import StorageBackend
15      from backend.storage.ref import StorageRef, StorageStat
16
17
18      class LocalStorage(StorageBackend):
19          def __init__(self, *, base_dir: str = "", **_ignored) -> None:
20              # Relative keys resolve under base_dir when set; absolute keys are used as-is.
21              self.base_dir = base_dir
22
23          def _path(self, ref: StorageRef) -> str:
24              p = ref.key
25              if self.base_dir and not os.path.isabs(p):
26                  return os.path.join(self.base_dir, p)
27              return p
28
29          def fetch_to_local(self, ref, dst=None, *, overwrite=False) -> str:
30              # The object is already on the local FS ? return its path (identity). ``dst`` is
31              # honoured only if it names a *different* location (a copy is then made).
32              src = self._path(ref)
33              if dst is None or os.path.abspath(str(dst)) == os.path.abspath(src):
34                  return src
35              dst = str(dst)
36              if os.path.exists(dst) and not overwrite:
37                  return dst
38              os.makedirs(os.path.dirname(dst) or ".", exist_ok=True)
39              shutil.copyfile(src, dst)
40              return dst
41
42          def put(self, local_path, ref, *, content_type=None) -> StorageRef:
43              dst = self._path(ref)
44              if os.path.abspath(str(local_path)) != os.path.abspath(dst):
45                  os.makedirs(os.path.dirname(dst) or ".", exist_ok=True)
46                  shutil.copyfile(str(local_path), dst)
```

```
47              return ref
48
49          def exists(self, ref) -> bool:
50              return os.path.exists(self._path(ref))
51
52          def stat(self, ref) -> StorageStat:
53              st = os.stat(self._path(ref))
54              return StorageStat(size=st.st_size,
modified=datetime.fromtimestamp(st.st_mtime))
55
56          def list(self, prefix) -> Iterator[StorageRef]:
57              base = self._path(prefix)
58              if os.path.isdir(base):
59                  for root, _dirs, files in os.walk(base):
60                      for fn in sorted(files):
61                          full = os.path.join(root, fn)
62                          yield StorageRef("local", "", full, f"local://{full}")
63              elif os.path.exists(base):
64                  yield StorageRef("local", "", base, f"local://{base}")
65
66          def delete(self, ref) -> None:
67              path = self._path(ref)
68              if os.path.exists(path):
69                  os.remove(path)
70
```

9. assistant (2026-06-20T09:01:13.112Z)

Now I can confirm the key mechanics. `LocalStorage.list` recurses (`os.walk`) and yields ALL files with no extension filter, in sorted order per directory. Let me verify the `parse_uri` `.name` property and check what `compute_video_id` and the runner do, plus confirm the labels side filters `.csv`.

10. user (2026-06-20T09:01:13.693Z)

```
1       """Storage value objects + URI parsing for the pluggable storage layer.
2
3       A ``StorageRef`` is a backend-agnostic pointer to one object. URIs:
4
5           <bare path>                -> local   (today's behaviour; e.g. C:/v/x.mp4,
/data/x.mp4)
6           local://<path>             -> local   (local:///abs or local://rel/p)
7           s3://<bucket>/<key>        -> s3       (AWS + R2 + DO Spaces + MinIO via
endpoint_url)
8           gcs://<bucket>/<key>       -> gcs      (gs:// accepted, normalized to gcs://)
9           gdrive://<folderId>/<path>  -> gdrive   (follow-up backend)
10
11      No cloud SDKs are imported here ? this module is pure stdlib so ``import
backend.storage``
12      stays cheap.
13      """
14      from __future__ import annotations
15
16      import re
17      from dataclasses import dataclass
18      from datetime import datetime
19      from typing import Optional
20
21      _SCHEME_RE = re.compile(r"^([A-Za-z][A-Za-z0-9+.\-]*)://")
22      _KNOWN = ("local", "s3", "gcs", "gdrive")
23
24
25      @dataclass(frozen=True)
26      class StorageStat:
27          """Lightweight object metadata (uniform across backends)."""
```

```
28          size: int
29          modified: Optional[datetime] = None
30          etag: Optional[str] = None
31          content_type: Optional[str] = None
32
33
34      @dataclass(frozen=True)
35      class StorageRef:
36          """A backend + a location. ``bucket`` is the S3/GCS bucket or the Drive folder-id;
37          for ``local`` it is empty and ``key`` carries the filesystem path."""
38          backend: str
39          bucket: str
40          key: str
41          uri: str
42
43          @property
44          def name(self) -> str:
45              return self.key.rstrip("/").rsplit("/", 1)[-1]
46
47
48      def parse_uri(uri: str) -> StorageRef:
49          """Parse a storage URI (or a bare filesystem path) into a ``StorageRef``.
50
51          A string with no ``scheme://`` prefix (including Windows ``C:\\...`` paths, which
52          have no ``//``) is treated as a local filesystem path ? preserving today's
behaviour.
53          """
54          s = str(uri)
55          m = _SCHEME_RE.match(s)
56          if not m:
57              return StorageRef("local", "", s, f"local://{s}")
58          scheme = m.group(1).lower()
59          rest = s[m.end():]
60          if scheme == "gs":
61              scheme = "gcs"
62          if scheme == "local":
63              return StorageRef("local", "", rest, s)
64          if scheme in _KNOWN:
65              bucket, _, key = rest.partition("/")
66              return StorageRef(scheme, bucket, key, s)
67          raise ValueError(f"unsupported storage URI scheme: {scheme}:// (in {s!r})")
68
```

11. assistant (2026-06-20T09:01:16.183Z)

Now let me check the S3 storage `list` to confirm cross-backend behavior, and verify the runner's `run_predict` to understand what happens when fed a non-video file.

12. user (2026-06-20T09:01:16.983Z)

```
1       """S3 (and S3-compatible) storage backend ? first-class object buckets.
2
3       One backend covers **AWS S3, Cloudflare R2, DigitalOcean Spaces, and MinIO** ? they
differ
4       only by ``endpoint_url`` (+ region / addressing style). Uses ``boto3``, imported lazily
so
5       ``import backend.storage`` never pulls it; install with the ``storage-s3`` extra.
6
7       Credentials are NEVER read from config ? only ``endpoint_url`` / ``region`` / addressing
8       style. boto3's default chain supplies creds (``AWS_ACCESS_KEY_ID`` /
``AWS_SECRET_ACCESS_KEY``
9       / ``AWS_SESSION_TOKEN`` env, instance role, or ``~/.aws``). A ``client`` may be injected
10      (tests pass a fake; production lets the constructor build one via boto3).
11      """
12      from __future__ import annotations
```

```

    import os
    from typing import Any, Iterator, Optional

    from backend.storage.base import StorageBackend
    from backend.storage.ref import StorageRef, StorageStat


    class S3Storage(StorageBackend):
        def __init__(self, *, client: Any = None, endpoint_url: Optional[str] = None,
                     region: Optional[str] = None, addressing_style: str = "auto",
                     signed_url_ttl: int = 3600, **_ignored) -> None:
            self._ttl = signed_url_ttl
            if client is not None:
                self._client = client
                return
            import boto3  # type: ignore  # lazy: only when an s3:// URI is actually used
            from botocore.config import Config  # type: ignore
            self._client = boto3.client(
                "s3",
                endpoint_url=endpoint_url or os.environ.get("AWS_ENDPOINT_URL_S3") or None,
                region_name=region,
                config=Config(s3={"addressing_style": addressing_style}),
            )

        def fetch_to_local(self, ref, dst=None, *, overwrite=False) -> str:
            if dst is None:
                raise ValueError("S3Storage.fetch_to_local requires a dst path")
            dst = str(dst)
            if os.path.exists(dst) and not overwrite:
                return dst
            os.makedirs(os.path.dirname(dst) or ".", exist_ok=True)
            self._client.download_file(ref.bucket, ref.key, dst)
            return dst

        def put(self, local_path, ref, *, content_type=None) -> StorageRef:
            extra = {"ContentType": content_type} if content_type else None
            self._client.upload_file(str(local_path), ref.bucket, ref.key, ExtraArgs=extra)
            return ref

        def exists(self, ref) -> bool:
            try:
                self._client.head_object(Bucket=ref.bucket, Key=ref.key)
                return True
            except Exception:  # noqa: BLE001 - any boto/client error means "not found /
unreachable"
                return False

        def stat(self, ref) -> StorageStat:
            h = self._client.head_object(Bucket=ref.bucket, Key=ref.key)
            return StorageStat(size=int(h.get("ContentLength", 0)),
modified=h.get("LastModified"),
                               etag=h.get("ETag"), content_type=h.get("ContentType"))

        def list(self, prefix) -> Iterator[StorageRef]:
            paginator = self._client.get_paginator("list_objects_v2")
            for page in paginator.paginate(Bucket=prefix.bucket, Prefix=prefix.key):
                for obj in page.get("Contents", []) or []:
                    key = obj["Key"]
                    yield StorageRef("s3", prefix.bucket, key,
f"s3://{prefix.bucket}/{key}")

        def delete(self, ref) -> None:
            self._client.delete_object(Bucket=ref.bucket, Key=ref.key)

```

```
75          def signed_url(self, ref, *, ttl=3600, method="GET") -> str:
76              op = "get_object" if method.upper() == "GET" else "put_object"
77              return self._client.generate_presigned_url(
78                  op, Params={"Bucket": ref.bucket, "Key": ref.key}, ExpiresIn=ttl or
self._ttl)
79
```

13. assistant (2026-06-20T09:01:20.389Z)

S3 `list` confirms: it returns ALL keys under the prefix with no extension filter, ordered as the bucket returns them (lexicographic for S3). Now let me examine the native runner's `run_predict` to understand the consequence of being handed a non-video file, and check the stub runner behavior.

14. user (2026-06-20T09:01:20.930Z)

```
1       """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3       The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4       "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
M3.5).
5       It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
6       Linux GPU box (GCP / a SaaS worker): it runs the **same** verified ``wasb_infer.py``
core
7       that the WSL adapter runs, so the trajectory CSV is byte-identical ? but WITHOUT the WSL
8       transport:
9
10        * NO ``WslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
11          ``wasb`` env's python is invoked **by path**), NO ``/mnt/c`` path translation
12          (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
13        * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps`` or ``cpu``; threaded
14          into ``wasb_infer.py --device`` (and so into the Hydra ``runner.device`` override).
15        * weights / repo / python / scratch come from the **environment** first
16          (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
17          ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
18
19      It emits the identical ``Frame,Visibility,X,Y`` CSV (same ``{stem}__{vid12}_wasb.csv``
20      name as the WSL runner) and reuses the model-agnostic ``parse_trajectory_csv`` /
21      ``trajectory_to_action_windows`` from ``tracknet_runner``, so the trajectory hybrids and
22      the windowing brain are unchanged. Selected via the existing ``detector_impl`` seam
23      (``trajectory_hybrid._resolve_runner`` ? ``_build_native_runner``);
``detector_impl="native"``.
24
25      **Acceptance gate (owner / GPU box):** byte-parity ? the SAME clip through the WSL
runner
26      (Windows) and this runner (Linux, ``device=cuda``) yields an identical trajectory CSV
and
27      identical candidate windows. That parity is a hardware check (a real GPU + the ``wasb``
28      env), so it is verified by the owner on the box ? not in this offline, subprocess-mocked
29      test suite.
30      """
31      import os
32      import shutil
33      import logging
34      import subprocess
35      from dataclasses import dataclass
36      from typing import Any, Dict, List, Optional, Tuple
37
38      from backend.pipeline.detectors.base import DetectorRunner
39      # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
40      from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
41
42      logger = logging.getLogger(__name__)
43
```

```
    # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
    # detectors/trackers/dataloaders + find configs/), exactly as the WSL adapter does.
    _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")


    @dataclass
    class NativeWasbConfig:
        """Resolved settings for a Linux-native WASB run.

        Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
        over ``indexing.wasb.*`` config (the box sets env; no config edits needed).
``device``
        defaults to ``cuda`` (the WSL byte-parity path).
        """
        repo_dir: str = "~/models/WASB-SBDT"   # native clone of nttcom/WASB-SBDT
        python_bin: str = "python"             # the `wasb` conda env python (invoked by
path, no activate)
        weights_path: str = ""                 # REQUIRED (env WASB_WEIGHTS or
indexing.wasb.weights_path)
        sport: str = "badminton"
        device: str = "cuda"                   # cuda | mps | cpu
        scratch_dir: str = ""                  # frame-extraction scratch (env); "" = next
to the output dir
        timeout_sec: int = 1800               # 30 min; a hung GPU call is killed (0 = no
timeout)
        keep_frames: bool = False              # retain the decoded frame cache after
success (debug only)
        stream_video: bool = False             # fast path: decode in-process, no PNG
extraction (bit-identical)

        @classmethod
        def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "NativeWasbConfig":
            w = (idx_cfg or {}).get("wasb", {}) or {}
            env = os.environ.get
            return cls(
                repo_dir=env("WASB_REPO_DIR") or w.get("repo_dir", "~/models/WASB-SBDT"),
                python_bin=env("WASB_PYTHON") or w.get("python_bin", "python"),
                weights_path=env("WASB_WEIGHTS") or w.get("weights_path", "") or "",
                sport=w.get("sport", "badminton"),
                device=env("WASB_DEVICE") or w.get("device", "cuda"),
                scratch_dir=env("WASB_SCRATCH") or env("RALLY_SCRATCH") or env("TMPDIR") or
"",
                timeout_sec=int(w.get("timeout_sec", 1800)),
                keep_frames=bool(w.get("keep_frames", False)),
                stream_video=bool(w.get("stream_video", False)),
            )


    class NativeWasbRunner(DetectorRunner):
        """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
docstring."""

        def __init__(self, cfg: NativeWasbConfig):
            self.cfg = cfg

        # --- low-level native exec (the single place tests patch)
-------------------------
        def _run(self, argv: List[str], cwd: Optional[str] = None) ->
subprocess.CompletedProcess:
            logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
            return subprocess.run(argv, cwd=cwd, capture_output=True, text=True,
                                  timeout=(self.cfg.timeout_sec or None))

        def _repo_src(self) -> str:
            return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
```

```
98
99          def _copy_infer_script(self) -> bool:
100             """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
configs/."""
101             repo_src = self._repo_src()
102             try:
103                 os.makedirs(repo_src, exist_ok=True)
104                 shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
105                 return True
106             except OSError as e:
107                 logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
108                 return False
109
110         def _rm_rf(self, path: Optional[str]) -> None:
111             """Best-effort recursive delete of a native path (post-success cleanup; never
raises)."""
112             if path:
113                 shutil.rmtree(path, ignore_errors=True)
114
115         def healthcheck(self) -> Tuple[bool, str]:
116             """Verify weights + repo present and the `wasb` python imports torch (and sees a
GPU on cuda)."""
117             if not self.cfg.weights_path:
118                 return False, ("WASB weights path is empty ? set WASB_WEIGHTS (or
indexing.wasb.weights_path).")
119             weights = os.path.expanduser(self.cfg.weights_path)
120             if not os.path.exists(weights):
121                 return False, f"WASB weights not found at {weights}"
122             repo_src = self._repo_src()
123             if not os.path.isdir(repo_src):
124                 return False, (f"WASB-SBDT repo src/ not found at {repo_src} ? clone
nttcom/WASB-SBDT "
125                                f"(set WASB_REPO_DIR / indexing.wasb.repo_dir).")
126             code = "import torch; print('torch', torch.__version__, 'cuda',
torch.cuda.is_available())"
127             try:
128                 res = self._run([self.cfg.python_bin, "-c", code])
129             except FileNotFoundError:
130                 return False, f"python binary not found: {self.cfg.python_bin!r} (set
WASB_PYTHON)."
131             except subprocess.TimeoutExpired:
132                 return False, "native torch healthcheck timed out."
133             out = (res.stdout or "").strip()
134             if res.returncode != 0:
135                 return False, f"`{self.cfg.python_bin}` torch import failed: {(res.stderr or
out).strip()}"
136             if self.cfg.device == "cuda" and "cuda True" not in out:
137                 return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
138             return True, out
139
140         def run_predict(self, video_win_path: str, output_win_dir: str,
141                         video_id: Optional[str] = None) -> Optional[str]:
142             """Run WASB natively on a video. Returns the path to the trajectory CSV, or
None.
143
144             Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
WSL
145             runner does, and the CSV is named ``{stem}__{vid12}_wasb.csv`` ? byte-for-byte
the
146             same downstream contract, so the parity comparison is apples-to-apples.
147             """
148             output_dir = os.path.abspath(output_win_dir)
149             os.makedirs(output_dir, exist_ok=True)
150             # Normalize for a cross-platform basename (tests pass Windows paths on Linux).
151             norm = str(video_win_path).replace("\\", "/")
```

```
            stem, _ext = os.path.splitext(os.path.basename(norm))
            # Content-derived key: same filename + different bytes -> different key (no
cache reuse).
            key = f"{stem}__{video_id[:12]}" if video_id else (stem or "wasb")
            out_csv = os.path.join(output_dir, f"{key}_wasb.csv")

            if not self._copy_infer_script():
                logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
                return None

            weights = os.path.expanduser(self.cfg.weights_path)
            argv = [self.cfg.python_bin, "wasb_infer.py",
                    "--video", os.path.abspath(str(video_win_path))]
            frames_dir: Optional[str] = None
            if self.cfg.stream_video:
                # Fast path: decode in-process, no PNG extraction. Output is bit-identical
to disk.
                argv.append("--stream-video")
            else:
                scratch = self.cfg.scratch_dir or output_dir
                frames_dir = os.path.join(scratch, f"{key}_frames")
                argv += ["--frames_out_dir", frames_dir]
            argv += ["--weights", weights,
                     "--sport", self.cfg.sport,
                     "--device", self.cfg.device,
                     "--out", out_csv]

            logger.info("Running native WASB inference (device=%s) on %s ...",
                        self.cfg.device, os.path.basename(norm))
            try:
                res = self._run(argv, cwd=self._repo_src())
            except FileNotFoundError:
                logger.error("python binary not found: %r (set WASB_PYTHON).",
self.cfg.python_bin)
                return None
            except subprocess.TimeoutExpired:
                logger.error("WASB native inference timed out after %ss.",
self.cfg.timeout_sec)
                return None
            if res.returncode != 0:
                logger.error("WASB native inference failed:\n%s", (res.stderr or res.stdout
or "")[-2000:])
                return None
            if not os.path.exists(out_csv):
                logger.error("WASB native finished but expected CSV not found at %s",
out_csv)
                return None
            logger.info("WASB trajectory CSV: %s", out_csv)

            # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
            # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
            # (Streaming never materializes frames, so there is nothing to drop on that
path.)
            if not self.cfg.keep_frames and frames_dir:
                logger.info("Cache hygiene: removing frame cache for %s "
                            "(set indexing.wasb.keep_frames=true to retain).", key)
                self._rm_rf(frames_dir)
            return out_csv

        # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
        parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
        trajectory_to_action_windows =
```

```
staticmethod(TrackNetRunner.trajectory_to_action_windows)
207
```

15. assistant (2026-06-20T09:01:27.604Z)

Now I understand the runner fully. When fed a non-video file (e.g. `v1.srt`), `wasb_infer.py` would be invoked with `--video v1.srt`. Let me confirm what would happen ? whether it fails loudly (returncode != 0 ? returns None ? "skipping") or silently succeeds. The key question for the finding's "silent corruption" severity. Let me check the labels-side filter and verify the realistic order of files.

Let me verify the exact claim about ordering. The finding says for `[v1.json, v1.mp4, v1.srt]` sorted, `video_by_stem['v1']` ends up as `v1.srt`. Since `os.walk` yields `sorted(files)` per directory, the iteration order within one dir is `v1.json, v1.mp4, v1.srt`, and the dict keeps the LAST write = `v1.srt`. That's confirmed mechanically.

Let me check whether sidecars actually land in `videos/` in this project's conventions, and whether the worker's infer path (not just train) has the same pattern.

16. user (2026-06-20T09:01:29.644Z)

```
98:    local_video = storage.fetch(args.video_uri, os.path.join(scratch, "in.mp4"),
config=storage_cfg)
224:    # Real-detector source: resolve the runner once + index the bucket's videos/ by stem.
232:        for vuri in storage.list_uris(f"{corpus}/videos/", config=storage_cfg):
245:                print(f"[worker] no video for {name!r} under {corpus}/videos/ ? skipping.")
319:    pt.add_argument("--corpus-uri", required=True, help="prefix containing labels/ (+
videos/ for "
326:                          "(needs videos/ in the corpus; honours detector_impl).")
```

17. user (2026-06-20T09:01:30.049Z)

```
backend\api\uploads.py:50:    exts = [e.lower().lstrip(".") for e in (getattr(sec,
"allowed_upload_extensions", None) or ["mp4", "mov"])]
backend\config\models.py:288:    allowed_upload_extensions: List[str] = ["mp4", "mov"]
docs\CODE_MAP.md:115:[Omitted long matching line]
```

18. assistant (2026-06-20T09:01:34.672Z)

So the `infer` path (line 98) uses a single explicit `args.video_uri` ? not affected. The vulnerability is specific to the `cmd_train` detector branch. The project does have an `allowed_upload_extensions` config defaulting to `["mp4", "mov"]`, confirming the codebase's own notion of valid video extensions, and confirming the proposed-fix reference.

Now let me check the test that was added for the worker to see whether it tests the videos/ listing, and whether any existing convention puts sidecars in videos/. Let me also confirm how a corpus's `videos/` is populated.

19. user (2026-06-20T09:01:35.106Z)

```
1       """Worker dispatcher: pull (storage) ? run pipeline ? push (storage).
2
3       `cmd_infer` is exercised end-to-end over the LOCAL storage backend with the validator +
4       segmenter mocked, so it needs no real video/GPU/network. The real S3 round-trip is
validated
5       on the droplet against MinIO.
6       """
7       import json
8
9       import pytest
10
11      from backend.config import load_config
12      from backend.worker import __main__ as wmain
13
```

```
14
15      # ---------------------------------------------------------------- pure helpers
16
17      @pytest.mark.parametrize("raw,expected", [
18          ("true", True), ("False", False), ("3", 3), ("2.5", 2.5), ("wasb_hybrid",
"wasb_hybrid"),
19      ])
20      def test_coerce(raw, expected):
21          assert wmain._coerce(raw) == expected and type(wmain._coerce(raw)) is type(expected)
22
23
24      def test_apply_overrides_on_typed_config():
25          cfg = _fresh_config()
26          wmain._apply_overrides(cfg, ["indexing.skip_ai_handoff=true",
"indexing.min_segment_duration=1.5",
27                                       "sport=tennis"])
28          assert cfg.indexing.skip_ai_handoff is True
29          assert cfg.indexing.min_segment_duration == 1.5
30          assert cfg.sport == "tennis"
31
32
33      def test_apply_overrides_rejects_bad_format():
34          cfg = _fresh_config()
35          with pytest.raises(ValueError):
36              wmain._apply_overrides(cfg, ["no_equals_sign"])
37
38
39      def test_write_intervals_csv(tmp_path):
40          segs = [{"start_time": 2.0, "end_time": 5.0, "shot_count": 4,
"shot_count_confidence": "LocalNoAI"},
41                  {"start_time": 9.0, "end_time": 12.5}]
42          out = tmp_path / "intervals.csv"
43          wmain._write_intervals_csv(segs, str(out))
44          lines = out.read_text().splitlines()
45          assert lines[0] == "start,end,shot_count,ending_reason"
46          assert lines[1].startswith("2.000,5.000,4,")
47          assert lines[2].startswith("9.000,12.500,,")
48
49
50      def test_write_report_json(tmp_path):
51          out = tmp_path / "report.json"
52          wmain._write_report_json("vid1", [{"start_time": 1.0, "end_time": 2.0}], "success",
str(out))
53          data = json.loads(out.read_text())
54          assert data["video_id"] == "vid1" and data["segment_count"] == 1
55          assert data["segments"][0] == {"start": 1.0, "end": 2.0}
56
57
58      def test_parser_infer_accepts_set_and_uris():
59          args = wmain.build_parser().parse_args([
60              "infer", "--video-uri", "s3://b/v.mp4", "--out-uri", "s3://b",
61              "--set", "indexing.skip_ai_handoff=true", "--set", "sport=tennis"])
62          assert args.cmd == "infer" and args.video_uri == "s3://b/v.mp4"
63          assert args.set == ["indexing.skip_ai_handoff=true", "sport=tennis"]
64
65
66      # ---------------------------------------------------------------- cmd_infer (local,
mocked)
67
68      class _OKResult:
69          passed = True
70          validator_name = "fake"
71          message = "ok"
72
73          def model_dump(self):
```

```
            return {"validator_name": "fake", "passed": True, "message": "ok"}


    class _FailResult(_OKResult):
        passed = False
        message = "too blurry"


    class _FakeSeg:
        def __init__(self, db, config):
            pass

        def process_video(self, video_id, path, max_frames=None):
            return [
                {"start_time": 2.0, "end_time": 5.0, "shot_count": 4,
"shot_count_confidence": "LocalNoAI"},
                {"start_time": 9.0, "end_time": 12.0, "shot_count": 6,
"shot_count_confidence": "LocalNoAI"},
            ]


    def _fresh_config():
        # Load defaults without touching any cwd config.json.
        import tempfile
        p = tempfile.NamedTemporaryFile("w", suffix=".json", delete=False)
        p.write("{}")
        p.close()
        return load_config(p.name)


    @pytest.fixture
    def mocked_pipeline(monkeypatch):
        monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vid123")
        monkeypatch.setattr(wmain.validator_registry, "run_all_with_context",
                            lambda path, cfg: ([_OKResult()], {}))
        monkeypatch.setattr(wmain.segmenter_registry, "get", lambda name: _FakeSeg)
        monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)


    def test_cmd_infer_local_roundtrip(tmp_path, mocked_pipeline):
        video = tmp_path / "in.mp4"
        video.write_bytes(b"FAKEVIDEO")
        cfg = tmp_path / "config.json"
        cfg.write_text("{}")
        out = tmp_path / "bucket"
        scratch = tmp_path / "scratch"

        rc = wmain.main([
            "infer", "--video-uri", str(video), "--out-uri", str(out),
            "--config", str(cfg), "--scratch", str(scratch),
            "--segmenter", "wasb_hybrid", "--set", "indexing.skip_ai_handoff=true",
        ])
        assert rc == 0
        intervals = out / "outputs" / "vid123" / "intervals.csv"
        report = out / "outputs" / "vid123" / "report.json"
        assert intervals.exists() and report.exists()
        rows = intervals.read_text().splitlines()
        assert len(rows) == 3                           # header + 2 rallies
        assert json.loads(report.read_text())["segment_count"] == 2


    def test_cmd_infer_validation_failure_returns_2(tmp_path, monkeypatch):
        monkeypatch.setattr(wmain, "compute_video_id", lambda p: "vidF")
        monkeypatch.setattr(wmain.validator_registry, "run_all_with_context",
                            lambda path, cfg: ([_FailResult()], {}))
```

```
        monkeypatch.setattr(wmain, "emit_indexer_report", lambda **k: None)
        video = tmp_path / "in.mp4"
        video.write_bytes(b"X")
        cfg = tmp_path / "config.json"
        cfg.write_text("{}")
        rc = wmain.main(["infer", "--video-uri", str(video), "--out-uri", str(tmp_path /
"o"),
                         "--config", str(cfg), "--scratch", str(tmp_path / "s")])
        assert rc == 2


    # --- cmd_train: pull labels -> synth trajectories -> manifest -> harness -> push bundle
---

    def test_cmd_train_orchestration(tmp_path, monkeypatch):
        import os
        import subprocess

        labels_dir = tmp_path / "corpus" / "labels"
        labels_dir.mkdir(parents=True)
        header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
        (labels_dir / "v1.rallies.csv").write_text(header +
"1,2.0,5.0,winner,badminton,\n2,9.0,12.0,let,badminton,\n")
        (labels_dir / "v2.rallies.csv").write_text(header + "1,3.0,6.0,winner,badminton,\n")
        out_dir = tmp_path / "out"
        scratch = tmp_path / "scratch"
        cfg = tmp_path / "config.json"
        cfg.write_text("{}")

        # Mock the harness subprocess: write a fake versioned bundle at --out/--version.
        def fake_run(cmd, *a, **k):
            out = cmd[cmd.index("--out") + 1]
            ver = cmd[cmd.index("--version") + 1]
            bdir = os.path.join(out, ver)
            os.makedirs(bdir, exist_ok=True)
            for fn in ("weights.json", "manifest.json"):
                with open(os.path.join(bdir, fn), "w") as f:
                    f.write("{}")
            return subprocess.CompletedProcess(cmd, 0)
        monkeypatch.setattr(subprocess, "run", fake_run)

        rc = wmain.main(["train", "--corpus-uri", str(tmp_path / "corpus"), "--out-uri",
str(out_dir),
                         "--version", "0.0.1", "--config", str(cfg), "--scratch",
str(scratch)])
        assert rc == 0
        # bundle pushed to <out-uri>/weights/<version>/
        assert (out_dir / "weights" / "0.0.1" / "weights.json").exists()
        assert (out_dir / "weights" / "0.0.1" / "manifest.json").exists()
        # the manifest fed to the harness enumerated both labeled videos with a trajectory
each
        specs = json.loads((scratch / "manifest.json").read_text())
        assert len(specs) == 2
        assert all("trajectory" in s and "labels" in s and "fps" in s for s in specs)


    def test_parser_train_detector_args():
        args = wmain.build_parser().parse_args([
            "train", "--corpus-uri", "gcs://b", "--out-uri", "gcs://b", "--version",
"1.0.0",
            "--trajectory-source", "detector", "--segmenter", "fusion_hybrid"])
        assert args.trajectory_source == "detector" and args.segmenter == "fusion_hybrid"


    def test_cmd_train_detector_source_runs_real_extract_via_stub(tmp_path, monkeypatch):
```

```
195         """--trajectory-source detector pulls videos + runs the resolved DetectorRunner.
With
196         RALLY_DETECTOR_IMPL=stub the *wiring* is exercised GPU-free (the stub IS the
runner)."""
197         import os
198         import subprocess
199
200         monkeypatch.setenv("RALLY_DETECTOR_IMPL", "stub")  # resolve a CPU stub runner (no
GPU/WSL)
201         corpus = tmp_path / "corpus"
202         labels_dir = corpus / "labels"
203         videos_dir = corpus / "videos"
204         labels_dir.mkdir(parents=True)
205         videos_dir.mkdir(parents=True)
206         header = "rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
207         (labels_dir / "v1.rallies.csv").write_text(header + "1,2.0,5.0,winner,badminton,\n")
208         (labels_dir / "v2.rallies.csv").write_text(header + "1,3.0,6.0,winner,badminton,\n")
209         (videos_dir / "v1.mp4").write_bytes(b"FAKEVIDEO1")
210         (videos_dir / "v2.mp4").write_bytes(b"FAKEVIDEO2")
211         out_dir = tmp_path / "out"
212         scratch = tmp_path / "scratch"
213         cfg = tmp_path / "config.json"
214         cfg.write_text("{}")
215
216         def fake_run(cmd, *a, **k):
217             out = cmd[cmd.index("--out") + 1]
218             ver = cmd[cmd.index("--version") + 1]
219             bdir = os.path.join(out, ver)
220             os.makedirs(bdir, exist_ok=True)
221             for fn in ("weights.json", "manifest.json"):
222                 with open(os.path.join(bdir, fn), "w") as f:
223                     f.write("{}")
224             return subprocess.CompletedProcess(cmd, 0)
225         monkeypatch.setattr(subprocess, "run", fake_run)
226
227         rc = wmain.main(["train", "--corpus-uri", str(corpus), "--out-uri", str(out_dir),
228                          "--version", "0.0.2", "--config", str(cfg), "--scratch",
str(scratch),
229                          "--trajectory-source", "detector", "--segmenter", "wasb_hybrid"])
230         assert rc == 0
231         specs = json.loads((scratch / "manifest.json").read_text())
232         assert len(specs) == 2
233         # Real-extract path: each spec points at an on-disk trajectory the runner actually
wrote.
234         assert all(os.path.exists(s["trajectory"]) for s in specs)
235         assert all(s["trajectory"].endswith("_stub.csv") for s in specs)  # the stub
runner's output
236
```

20. assistant (2026-06-20T09:01:38.638Z)

The test only exercises a clean `videos/` directory with exactly `v1.mp4` and `v2.mp4` ? no sidecars. So the bug is not covered by tests and the happy path masks it.

Now let me examine the stub_runner to understand the "blindly succeeds on garbage bytes" half of the claim, which affects severity (silent corruption vs. loud skip).

21. user (2026-06-20T09:01:39.211Z)

```
1       """CPU stub shuttle-trajectory runner ? "mock a GPU with a CPU".
2
3       A ``DetectorRunner`` that produces a deterministic ``Frame,Visibility,X,Y`` trajectory
4       CSV with **no GPU, no WSL, no model weights** ? so the whole trajectory-hybrid pipeline
5       (``wasb_hybrid`` / ``tracknet_hybrid``) can run end-to-end on a CPU-only box (CI, a CPU
6       container, a laptop) for regression testing. The person/yolo path already CPU-falls-back
```

```
    (``use_gpu=False``); this closes the last GPU/WSL-only gap so the *entire* pipeline is
    runnable ? and regression-gateable ? without a GPU.

    Two modes:
      * **replay** ? emit an existing ``Frame,Visibility,X,Y`` CSV verbatim (drive the
        pipeline with a committed/golden trajectory; deterministic + cross-OS).
      * **synth** (default) ? generate a deterministic synthetic trajectory of
        ``num_rallies`` in-flight bursts (visible + fast-moving frames) separated by dead
        gaps, so windowing yields a stable, known number of candidate windows.

    This is the CPU seam the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
    "Optional Linux-native path") and the **seed of the future ``LinuxNativeRunner``**
    (M3.5, ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``). It is selected via
    ``indexing.detector_impl="stub"`` or the ``RALLY_DETECTOR_IMPL=stub`` env var; the real
    WSL runners stay the default, so production behaviour is unchanged.
    """
    import os
    import csv as _csv
    import shutil
    import logging
    from dataclasses import dataclass
    from typing import Any, Dict, List, Optional, Tuple

    from backend.pipeline.detectors.base import DetectorRunner
    # Reuse the model-agnostic CSV parsing + windowing (no torch/tf pulled here).
    from backend.pipeline.detectors.tracknet_runner import TrackNetRunner

    logger = logging.getLogger(__name__)


    @dataclass
    class StubConfig:
        """Settings for the CPU stub runner (built from ``config.json`` ->
``indexing.stub``)."""
        fps: float = 30.0
        frame_width: float = 1280.0
        num_rallies: int = 2
        rally_seconds: float = 4.0
        gap_seconds: float = 3.0          # dead gap between rallies (> merge_gap so
windows don't merge)
        lead_seconds: float = 2.0         # dead lead-in
        step_px: float = 12.0             # per-frame shuttle displacement during a rally
(>> velocity_thresh)
        replay_csv: Optional[str] = None  # if set + exists, emit this CSV verbatim instead
of synthesising

        @classmethod
        def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "StubConfig":
            s = (idx_cfg or {}).get("stub", {}) or {}
            return cls(
                fps=float(s.get("fps", 30.0)),
                frame_width=float(s.get("frame_width", 1280.0)),
                num_rallies=int(s.get("num_rallies", 2)),
                rally_seconds=float(s.get("rally_seconds", 4.0)),
                gap_seconds=float(s.get("gap_seconds", 3.0)),
                lead_seconds=float(s.get("lead_seconds", 2.0)),
                step_px=float(s.get("step_px", 12.0)),
                replay_csv=s.get("replay_csv"),
            )


    class StubDetectorRunner(DetectorRunner):
        """Deterministic CPU trajectory runner (no GPU/WSL). See the module docstring."""

        def __init__(self, cfg: Optional[StubConfig] = None):
```

```
68              self.cfg = cfg or StubConfig()
69
70          def healthcheck(self) -> Tuple[bool, str]:
71              return True, "stub CPU detector (no GPU/WSL)"
72
73          def _rows(self) -> List[Tuple[int, int, float, float]]:
74              """Deterministic ``(frame, visible, x, y)`` rows: ``num_rallies`` in-flight
bursts
75              separated by dead gaps. During a rally the shuttle oscillates ``step_px``/frame
?
76              well above the default 0.02 velocity threshold ? so every rally frame is
"active"
77              and each burst becomes exactly one candidate window."""
78              cfg = self.cfg
79              rows: List[Tuple[int, int, float, float]] = []
80              frame = 0
81              y = cfg.frame_width / 4.0
82              base_x = cfg.frame_width / 2.0
83              for _ in range(int(cfg.lead_seconds * cfg.fps)):
84                  rows.append((frame, 0, 0.0, 0.0))
85                  frame += 1
86              for r in range(cfg.num_rallies):
87                  for i in range(int(cfg.rally_seconds * cfg.fps)):
88                      x = base_x + (cfg.step_px if i % 2 else -cfg.step_px)
89                      rows.append((frame, 1, x, y))
90                      frame += 1
91                  if r != cfg.num_rallies - 1:
92                      for _ in range(int(cfg.gap_seconds * cfg.fps)):
93                          rows.append((frame, 0, 0.0, 0.0))
94                          frame += 1
95              return rows
96
97          def run_predict(self, video_win_path: str, output_win_dir: str,
98                          video_id: Optional[str] = None) -> Optional[str]:
99              """Write a deterministic trajectory CSV and return its path (never touches the
GPU/WSL)."""
100             os.makedirs(output_win_dir, exist_ok=True)
101             stem = os.path.splitext(os.path.basename(str(video_win_path).replace("\\",
"/")))[0]
102             key = f"{stem}__{video_id[:12]}" if video_id else (stem or "stub")
103             out_csv = os.path.join(output_win_dir, f"{key}_stub.csv")
104
105             if self.cfg.replay_csv and os.path.exists(self.cfg.replay_csv):
106                 shutil.copyfile(self.cfg.replay_csv, out_csv)
107                 logger.info("Stub runner: replayed trajectory %s -> %s",
self.cfg.replay_csv, out_csv)
108                 return out_csv
109
110             rows = self._rows()
111             with open(out_csv, "w", newline="") as f:
112                 w = _csv.writer(f)
113                 w.writerow(["Frame", "Visibility", "X", "Y"])
114                 for frame, vis, x, y in rows:
115                     w.writerow([frame, vis, f"{x:.3f}", f"{y:.3f}"])
116             logger.info("Stub runner: wrote %d synthetic trajectory rows -> %s", len(rows),
out_csv)
117             return out_csv
118
119         # Model-agnostic CSV parsing + windowing ? identical to the real runners.
120         parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
121         trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)
122
123
124     def synth_trajectory_from_labels(labels_csv: str, out_csv: str, *, fps: float = 30.0,
```

```
125                                     frame_width: float = 1280.0, tail_pad: float = 2.0,
126                                     step_px: float = 12.0) -> str:
127         """CPU "mock GPU" for TRAINING: write a deterministic Frame,Visibility,X,Y
trajectory whose
128         in-flight (visible + oscillating) frames cover each labeled rally window and whose
dead frames
129         cover the gaps ? so the synthetic trajectory *encodes* the rally labels.
130
131         This lets the Gen-0 harness train + leave-one-video-out eval end-to-end on coherent
data with
132         no GPU/WSL, validating the whole training pipeline before the real WASB detector
lands. (The
133         learned model is meaningful for plumbing, not for production ? the trajectory is
synthetic.)
134         """
135         rallies: List[Tuple[float, float]] = []
136         with open(labels_csv, newline="", encoding="utf-8") as f:
137             for row in _csv.DictReader(f):
138                 try:
139                     rallies.append((float(row["start_time"]), float(row["end_time"])))
140                 except (KeyError, ValueError):
141                     continue
142         duration = (max(e for _, e in rallies) + tail_pad) if rallies else 10.0
143         nframes = max(1, int(duration * fps))
144         base_x = frame_width / 2.0
145         y = frame_width / 4.0
146         os.makedirs(os.path.dirname(out_csv) or ".", exist_ok=True)
147         with open(out_csv, "w", newline="", encoding="utf-8") as f:
148             w = _csv.writer(f)
149             w.writerow(["Frame", "Visibility", "X", "Y"])
150             for i in range(nframes):
151                 t = i / fps
152                 if any(s <= t < e for s, e in rallies):
153                     x = base_x + (step_px if i % 2 else -step_px)
154                     w.writerow([i, 1, f"{x:.3f}", f"{y:.3f}"])
155                 else:
156                     w.writerow([i, 0, "0.000", "0.000"])
157         return out_csv
158
```

22. assistant (2026-06-20T09:01:49.155Z)

This is a critical confirmation of the severity nuance. The `StubDetectorRunner.run_predict` (line 97-117) **completely ignores the video bytes**. It writes a deterministic synthetic trajectory regardless of what `video_win_path` points at. So when `RALLY_DETECTOR_IMPL=stub` is set and a sidecar shadows the real video:
- `compute_video_id` hashes the sidecar (`v1.srt`) ? wrong video_id (wrong cache key)
- The stub "succeeds" ? writes a trajectory keyed to the wrong video_id ? **pollutes the training manifest silently** (exactly as the finding states).

For the native runner on a real GPU box: `wasb_infer.py --video v1.srt` would almost certainly fail (ffmpeg/decoder can't decode a subtitle file) ? returncode != 0 ? `run_predict` returns None ? "detector produced no trajectory ... skipping" ? a real video is silently skipped (wrong-skip). Either path is a real defect.

Now let me confirm the finding's exact line-number claims and the precise labels-side `.csv` filter contrast. The finding cites line 217 (labels filter) and lines 232-234 / 244 (videos build + lookup). My read confirms:
- Line 217-218: labels filtered to `.csv`.
- Lines 232-234: `video_by_stem[stem] = vuri` with NO extension filter.
- Line 243-246: lookup + skip.

Let me also double check `compute_video_id` is the worker-imported one and there isn't a guard I'm missing in `_resolve_train_detector` or elsewhere that filters videos.

23. user (2026-06-20T09:01:50.313Z)

```
1       """Worker dispatcher ? pull from storage ? run the pipeline ? push outputs back.
2
3           python -m backend.worker infer --video-uri <uri> --out-uri <prefix> \
4               [--segmenter wasb_hybrid] [--config config.json] [--set k=v ...] [--scratch DIR]
5           python -m backend.worker train ...        # Gen-0 train-from-bucket (follow-up)
6
7       The URI scheme selects the storage backend via ``backend.storage`` (a bare path /
``local://``
8       is an identity passthrough; ``s3://`` / ``gcs://`` is an object store), so existing
local runs
9       are byte-for-byte unchanged. Combine with ``RALLY_DETECTOR_IMPL=stub`` and
10      ``--set indexing.skip_ai_handoff=true`` for a **secret-free, GPU-free** run (no Gemini
key) ?
11      the path validated end-to-end on DigitalOcean (``docs/SETUP_NEW_MACHINE.md``).
12
13      This is a thin orchestrator: it composes the SAME building blocks the ``backend.main``
CLI uses
14      (validator registry ? segmenter ? report), never a forked pipeline.
15      """
16      from __future__ import annotations
17
18      import argparse
19      import csv
20      import json
21      import os
22      import sys
23      import tempfile
24      from typing import Any, List
25
26      from backend import storage
27      from backend.config import load_config
28      from backend.infrastructure.database import Database
29      from backend.pipeline.segmenters.base import segmenter_registry
30      from backend.pipeline.validators.base import registry as validator_registry
31      from backend.reporting import emit_indexer_report
32      from backend.results import Failure
33      from backend.utils.hashing import compute_video_id
34
35
36      def _coerce(v: str) -> Any:
37          """Coerce a ``--set`` string value to bool/int/float where it round-trips, else
str."""
38          low = v.lower()
39          if low in ("true", "false"):
40              return low == "true"
```

24. assistant (2026-06-20T09:02:07.967Z)

I have everything needed to render a verdict. Let me finalize my assessment.

The finding is mechanically accurate and confirmed against the working tree:

1. **No extension filter on videos** (lines 232-234): `for vuri in storage.list_uris(f"{corpus}/videos/", ...)` then `video_by_stem[os.path.splitext(vname)[0]] = vuri` ? confirmed, no filtering by extension.

2. **Asymmetry with labels** (lines 217-218): labels ARE filtered to `.csv` via `if u.lower().endswith(".csv")`. The contrast is real and the videos side lacks the equivalent guard.

3. **Last-write-wins on stem collision**: confirmed by `LocalStorage.list` (`os.walk` + `sorted(files)`) yielding `v1.json, v1.mp4, v1.srt` in order, so the dict ends with `v1.srt`. S3 `list` similarly returns all keys lexicographically with no filter.

4. **Realistic silent corruption**: With `RALLY_DETECTOR_IMPL=stub` (an explicitly documented/supported wiring mode, used in `test_cmd_train_detector_source_runs_real_extract_via_stub`), `StubDetectorRunner.run_predict` ignores the video bytes entirely and "succeeds" on any path ? so a shadowing `v1.srt` produces a trajectory keyed to the wrong `compute_video_id` and pollutes the training manifest silently. On the native GPU path, `wasb_infer.py --video v1.srt` fails ? trajectory None ? the real video is silently skipped (wrong-skip). Both are real failure modes.

A couple of honest counterpoints to weigh on severity:
- This is a NEW code path (`--trajectory-source detector`), default source stays `synth`, so it is opt-in and not yet on the default e2e.
- It requires sidecars to actually be placed in `videos/` ? plausible per the project's bucket conventions but not proven to currently exist in the corpus.
- The native runner half degrades to a loud-ish skip (logged), not silent corruption; only the stub half is truly silent. That makes "silent corruption" partly true (stub) and partly an overstatement (native logs a skip).

Net: a real, confirmed defect with a correct proposed fix. Severity "major" is slightly high given it is an opt-in, not-yet-default path that on the production (native) runner degrades to a logged skip rather than silent corruption; "minor" understates the stub-path silent manifest pollution. I'll keep it as a legitimate finding but corrected to between ? leaning to keeping major is defensible, but the most honest call is minor-to-major. I'll mark it real with corrected severity "minor" because the truly-silent corruption requires the stub mode, and the default/native path fails loudly-ish; the data-correctness risk is real but gated behind an opt-in flag and atypical bucket contents.

25. user (2026-06-20T09:02:19.314Z)

Structured output provided successfully

26. assistant (2026-06-20T09:02:23.614Z)

Verdict: real (corrected to minor). The finding is mechanically accurate ? confirmed in the working tree across `backend/worker/__main__.py:217,232-234,243-246`, `backend/storage/local.py:56-64`, and `backend/storage/s3.py:65-70`. I downgraded severity from major to minor because the truly-silent corruption only occurs in the opt-in stub mode (`backend/pipeline/detectors/stub_runner.py:97-117` ignores video bytes), while the production native path degrades to a logged skip; the path is also default-OFF (`synth`) and requires sidecars in `videos/`. The proposed fix is correct.