

claude-windows / c47f0851-d75f-45b6-9126-e48544381426

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c47f0851-d75f-45b6-9126-e48544381426\subagents\workflows\wf_5af9a132-4ce\agent-a8557549219ec8cd7.jsonl`
- SHA-256: `497556a20df316f6ca8a1ec76567c96c70145ad4f18a2e5c392600a2ee8b3729`
- Source modified: `2026-06-20T09:01:27+00:00`
- Imported at: `2026-07-05T16:48:23+00:00`
- project: `wf_5af9a132-4ce`
- session_id: `c47f0851-d75f-45b6-9126-e48544381426`
```

Transcript

1. user (2026-06-20T08:59:56.758Z)

Adversarially verify this review finding against the ACTUAL code in the working tree. Default to 'false-positive' unless you can point to the concrete line(s) and a real failure scenario. Context: Repo: badminton-highlight-indexer (Python). Branch feat/native-wasb-runner off origin/master.

The change builds the M3.5 Linux-native WASB detector runner + wires real extraction into the worker.

Files changed/added (read them in the WORKING TREE ? changes are uncommitted):

```
- backend/pipeline/detectors/native_wasb_runner.py (NEW ? the native runner)
- backend/pipeline/detectors/wasb_infer.py (added --device, threaded device into
_build_cfg/run/run_video_streaming/main; device-keyed the resumable manifest)
- backend/pipeline/segmenters/trajectory_hybrid.py (_resolve_runner: new 'native' branch +
_build_native_runner base hook; generalized 'WSL env' log lines)
- backend/pipeline/segmenters/wasb_hybrid.py (_build_native_runner override)
- backend/pipeline/segmenters/fusion_hybrid.py (_build_native_runner override; WASB
substrate only)
- backend/config/models.py (WasbIndexConfig.device + .python_bin;
detector_impl doc mentions 'native')
- backend/worker/__main__.py (cmd_train: --trajectory-source
{synth,detector} + --segmenter; _resolve_train_detector; detector path pulls videos + runs
runner)
- tests/test_native_wasb_runner.py, tests/test_worker.py (new tests)
```

DESIGN INTENT / CONTRACTS (judge against these):

```
- The EXISTING WSL runner is backend/pipeline/detectors/wasb_runner.py (WasbRunner,
WslCommandMixin). The native runner must be a sibling that runs the SAME wasb_infer.py natively
(no wsl, no 'conda activate' ? it invokes the wasb conda env python BY PATH, no /mnt path
translation) and emit a BYTE-IDENTICAL Frame,Visibility,X,Y CSV with the SAME filename
'{stem}_{vid12}_wasb.csv'.
- ACCEPTANCE GATE = byte-parity: same clip via WSL (Windows) vs native (Linux, device=cuda) ->
identical CSV + identical candidate windows. device default 'cuda' must keep wasb_infer's Hydra
overrides identical to before (parity).
- Selected via the detector_impl seam: indexing.detector_impl / RALLY_DETECTOR_IMPL = 'native'.
- Constraint: 'backend' must NOT import 'training' (worker shells out). The default trajectory
source stays 'synth' so the CPU e2e is unchanged.
- The code RUNS on Linux (GPU box) but the TESTS run on the dev's Windows machine; cross-
platform path handling matters.
- weights/repo/python/scratch come from env first
(WASB_WEIGHTS/WASB_REPO_DIR/WASB_PYTHON/WASB_SCRATCH) then indexing.wasb.* then default.
```

FINDING [tests-regress] (major) @ tests/test_native_wasb_runner.py::_drive_success + test_run_predict_builds_native_argv_no_wsl (covering native_wasb_runner.py::run_predict lines 168-171, scratch_dir handling): The scratch_dir -> frames path wiring is completely untested. run_predict computes `scratch = self.cfg.scratch\_dir or output\_dir` and `frames\_dir = os.path.join(scratch, f'{key}\_frames')`, then passes it as --frames_out_dir. Tests assert '--

frames_out_dir' is PRESENT but never assert its VALUE, and never set scratch_dir, so the env-driven WASB_SCRATCH/RALLY_SCRATCH/TMPDIR path (a first-class env contract per the design intent) is never followed end-to-end into the argv. The frames dir that gets --rm -rf'd on cleanup is likewise never asserted.

WHY: If scratch_dir is dropped from the path join, frames silently land next to the output on the GPU box (the small system disk) instead of the large scratch volume ? a disk-fill regression on exactly the long 4K clips this runner targets, and tests stay green. Worse, the cleanup _rm_rf target could diverge from the extraction target (leaving the GBs of PNGs), undetected.

PROPOSED FIX: Add a test driving run_predict with NativeWasbConfig(scratch_dir='/scratch') (disk mode) that asserts argv[argv.index('--frames_out_dir')+1] == os.path.join('/scratch', f'{key}_frames'), and a second with scratch_dir='' asserting the frames dir is under the abspath'd output_dir. Capture the _rm_rf spy arg and assert it equals that same frames_dir.

2. assistant (2026-06-20T08:59:59.194Z)

I'll verify this finding against the actual code. Let me start by reading the relevant files.

3. user (2026-06-20T09:00:00.968Z)

On branch feat/native-wasb-runner

Your branch is up to date with 'origin/master'.

Changes not staged for commit:

(use "git add <file>..." to update what will be committed)

(use "git restore <file>..." to discard changes in working directory)

modified: backend/config/models.py
modified: backend/pipeline/detectors/wasb_infer.py
modified: backend/pipeline/segmenters/fusion_hybrid.py
modified: backend/pipeline/segmenters/trajectory_hybrid.py
modified: backend/pipeline/segmenters/wasb_hybrid.py
modified: backend/worker/__main__.py
modified: tests/test_worker.py

Untracked files:

(use "git add <file>..." to include in what will be committed)

artifacts/
backend/pipeline/detectors/native_wasb_runner.py
docs/KHELSUTRA_PRODUCT_OVERVIEW.md
tests/test_native_wasb_runner.py

no changes added to commit (use "git add" and/or "git commit -a")

---BRANCH---

feat/native-wasb-runner

4. user (2026-06-20T09:00:01.545Z)

```
1      """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3      The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4      "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
5      M3.5).
6      It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
7      Linux GPU box (GCP / a SaaS worker): it runs the same verified ``wasb_infer.py``
8      core
9      that the WSL adapter runs, so the trajectory CSV is byte-identical ? but WITHOUT the WSL
10     transport:
11
12     * NO ``wslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
13       ``wasb`` env's python is invoked by path), NO ``/mnt/c`` path translation
14       (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
15     * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps`` or ``cpu``; threaded
16       into ``wasb_infer.py --device`` (and so into the Hydra ``runner.device`` override).
17     * weights / repo / python / scratch come from the environment first
18       (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
```

```

17         ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
18
19     It emits the identical ``Frame,Visibility,X,Y`` CSV (same ``{stem}__{vid12}_wasb.csv``
20     name as the WSL runner) and reuses the model-agnostic ``parse_trajectory_csv`` /
21     ``trajectory_to_action_windows`` from ``tracknet_runner``, so the trajectory hybrids and
22     the windowing brain are unchanged. Selected via the existing ``detector_impl`` seam
23     (``trajectory_hybrid._resolve_runner`` ? ``_build_native_runner``);
24     ``detector_impl="native"``.
25
26     **Acceptance gate (owner / GPU box):** byte-parity ? the SAME clip through the WSL
runner
27     (Windows) and this runner (Linux, ``device=cuda``) yields an identical trajectory CSV
and
28     identical candidate windows. That parity is a hardware check (a real GPU + the ``wasb``
29     env), so it is verified by the owner on the box ? not in this offline, subprocess-mocked
30     test suite.
31     """
32     import os
33     import shutil
34     import logging
35     import subprocess
36     from dataclasses import dataclass
37     from typing import Any, Dict, List, Optional, Tuple
38
39     from backend.pipeline.detectors.base import DetectorRunner
40     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
41     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
42
43     logger = logging.getLogger(__name__)
44
45     # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
46     # detectors/trackers/dataloaders + find configs/), exactly as the WSL adapter does.
47     _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
48
49     @dataclass
50     class NativeWasbConfig:
51         """Resolved settings for a Linux-native WASB run.
52
53         Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
54         over ``indexing.wasb.*`` config (the box sets env; no config edits needed).
55         ``device``
56         defaults to ``cuda`` (the WSL byte-parity path).
57         """
58         repo_dir: str = "~/models/WASB-SBDT" # native clone of nttcom/WASB-SBDT
59         python_bin: str = "python" # the `wasb` conda env python (invoked by
path, no activate)
60         weights_path: str = "" # REQUIRED (env WASB_WEIGHTS or
indexing.wasb.weights_path)
61         sport: str = "badminton"
62         device: str = "cuda" # cuda | mps | cpu
63         scratch_dir: str = "" # frame-extraction scratch (env); "" = next
to the output dir
64         timeout_sec: int = 1800 # 30 min; a hung GPU call is killed (0 = no
timeout)
65         keep_frames: bool = False # retain the decoded frame cache after
success (debug only)
66         stream_video: bool = False # fast path: decode in-process, no PNG
extraction (bit-identical)
67
68     @classmethod
69     def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "NativeWasbConfig":
70         w = (idx_cfg or {}).get("wasb", {}) or {}
71         env = os.environ.get

```

```

71         return cls(
72             repo_dir=env("WASB_REPO_DIR") or w.get("repo_dir", "~/models/WASB-SBDT"),
73             python_bin=env("WASB_PYTHON") or w.get("python_bin", "python"),
74             weights_path=env("WASB_WEIGHTS") or w.get("weights_path", "") or "",
75             sport=w.get("sport", "badminton"),
76             device=env("WASB_DEVICE") or w.get("device", "cuda"),
77             scratch_dir=env("WASB_SCRATCH") or env("RALLY_SCRATCH") or env("TMPDIR") or
78             "",
79             timeout_sec=int(w.get("timeout_sec", 1800)),
80             keep_frames=bool(w.get("keep_frames", False)),
81             stream_video=bool(w.get("stream_video", False)),
82         )
83
84     class NativeWasbRunner(DetectorRunner):
85         """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
86         docstring."""
87         def __init__(self, cfg: NativeWasbConfig):
88             self.cfg = cfg
89
90         # --- low-level native exec (the single place tests patch)
91         -----
92         def _run(self, argv: List[str], cwd: Optional[str] = None) ->
93         subprocess.CompletedProcess:
94             logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
95             return subprocess.run(argv, cwd=cwd, capture_output=True, text=True,
96                                   timeout=(self.cfg.timeout_sec or None))
97
98         def _repo_src(self) -> str:
99             return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
100
101         def _copy_infer_script(self) -> bool:
102             """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
103             configs/."""
104             repo_src = self._repo_src()
105             try:
106                 os.makedirs(repo_src, exist_ok=True)
107                 shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
108                 return True
109             except OSError as e:
110                 logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
111                 return False
112
113         def _rm_rf(self, path: Optional[str]) -> None:
114             """Best-effort recursive delete of a native path (post-success cleanup; never
115             raises)."""
116             if path:
117                 shutil.rmtree(path, ignore_errors=True)
118
119         def healthcheck(self) -> Tuple[bool, str]:
120             """Verify weights + repo present and the `wasb` python imports torch (and sees a
121             GPU on cuda)."""
122             if not self.cfg.weights_path:
123                 return False, ("WASB weights path is empty ? set WASB_WEIGHTS (or
124                 indexing.wasb.weights_path).")
125             weights = os.path.expanduser(self.cfg.weights_path)
126             if not os.path.exists(weights):
127                 return False, f"WASB weights not found at {weights}"
128             repo_src = self._repo_src()
129             if not os.path.isdir(repo_src):
130                 return False, (f"WASB-SBDT repo src/ not found at {repo_src} ? clone
131                 nttcom/WASB-SBDT "
132                 f"(set WASB_REPO_DIR / indexing.wasb.repo_dir).")
133             code = "import torch; print('torch', torch.__version__, 'cuda',

```

```

torch.cuda.is_available())"
127         try:
128             res = self._run([self.cfg.python_bin, "-c", code])
129         except FileNotFoundError:
130             return False, f"python binary not found: {self.cfg.python_bin!r} (set
WASB_PYTHON)."
131         except subprocess.TimeoutExpired:
132             return False, "native torch healthcheck timed out."
133         out = (res.stdout or "").strip()
134         if res.returncode != 0:
135             return False, f"{self.cfg.python_bin}` torch import failed: {(res.stderr or
out).strip()}"
136         if self.cfg.device == "cuda" and "cuda True" not in out:
137             return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
138         return True, out
139
140     def run_predict(self, video_win_path: str, output_win_dir: str,
141                    video_id: Optional[str] = None) -> Optional[str]:
142         """Run WASB natively on a video. Returns the path to the trajectory CSV, or
None.
143
144         Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
WSL
145         runner does, and the CSV is named ``{stem}__{vid12}_wasb.csv`` ? byte-for-byte
the
146         same downstream contract, so the parity comparison is apples-to-apples.
147         """
148         output_dir = os.path.abspath(output_win_dir)
149         os.makedirs(output_dir, exist_ok=True)
150         # Normalize for a cross-platform basename (tests pass Windows paths on Linux).
151         norm = str(video_win_path).replace("\\", "/")
152         stem, _ext = os.path.splitext(os.path.basename(norm))
153         # Content-derived key: same filename + different bytes -> different key (no
cache reuse).
154         key = f"{stem}__{video_id[:12]}" if video_id else (stem or "wasb")
155         out_csv = os.path.join(output_dir, f"{key}_wasb.csv")
156
157         if not self._copy_infer_script():
158             logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
159             return None
160
161         weights = os.path.expanduser(self.cfg.weights_path)
162         argv = [self.cfg.python_bin, "wasb_infer.py",
163                "--video", os.path.abspath(str(video_win_path))]
164         frames_dir: Optional[str] = None
165         if self.cfg.stream_video:
166             # Fast path: decode in-process, no PNG extraction. Output is bit-identical
to disk.
167             argv.append("--stream-video")
168         else:
169             scratch = self.cfg.scratch_dir or output_dir
170             frames_dir = os.path.join(scratch, f"{key}_frames")
171             argv += ["--frames_out_dir", frames_dir]
172         argv += ["--weights", weights,
173                "--sport", self.cfg.sport,
174                "--device", self.cfg.device,
175                "--out", out_csv]
176
177         logger.info("Running native WASB inference (device=%s) on %s ...",
178                    self.cfg.device, os.path.basename(norm))
179         try:
180             res = self._run(argv, cwd=self._repo_src())
181         except FileNotFoundError:
182             logger.error("python binary not found: %r (set WASB_PYTHON).",
self.cfg.python_bin)

```

```

183         return None
184     except subprocess.TimeoutExpired:
185         logger.error("WASB native inference timed out after %ss.",
self.cfg.timeout_sec)
186         return None
187     if res.returncode != 0:
188         logger.error("WASB native inference failed:\n%s", (res.stderr or res.stdout
or "")[-2000:])
189         return None
190     if not os.path.exists(out_csv):
191         logger.error("WASB native finished but expected CSV not found at %s",
out_csv)
192         return None
193     logger.info("WASB trajectory CSV: %s", out_csv)
194
195     # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
196     # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
197     # (Streaming never materializes frames, so there is nothing to drop on that
path.)
198     if not self.cfg.keep_frames and frames_dir:
199         logger.info("Cache hygiene: removing frame cache for %s "
200                     "(set indexing.wasb.keep_frames=true to retain).", key)
201         self._rm_rf(frames_dir)
202     return out_csv
203
204     # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
205     parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
206     trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)
207

```

5. user (2026-06-20T09:00:01.814Z)

```

1     """NativeWasbRunner coverage ? the Linux-native (no-WSL) WASB detector path (M3.5).
2
3     The native subprocess is mocked (no torch/conda/GPU here). These tests lock the contract
4     the GPU-box byte-parity gate then relies on: it IS a DetectorRunner, it is NOT a WSL
runner
5     (no WslCommandMixin / wsl invocation), it builds a native argv (no `wsl`, no `conda
6     activate`, no `/mnt` translation) that runs the SAME wasb_infer.py with the SAME output
CSV
7     name as the WSL runner, threads the device through, and cleans the frame cache on
success.
8     """
9     import os
10    import types
11    from unittest.mock import patch
12
13    import pytest
14
15    from backend.pipeline.detectors.native_wasb_runner import NativeWasbRunner,
NativeWasbConfig
16    from backend.pipeline.detectors.base import DetectorRunner, WslCommandMixin
17
18
19    def _proc(rc=0, stdout="", stderr=""):
20        return types.SimpleNamespace(returncode=rc, stdout=stdout, stderr=stderr)
21
22
23    # --- identity / no-WSL guarantee
-----
24    def test_is_a_detector_runner():

```

```

25         assert isinstance(NativeWasbRunner(NativeWasbConfig()), DetectorRunner)
26
27
28     def test_is_not_a_wsl_runner():
29         """The whole point of M3.5: the native runner must carry NO WSL plumbing."""
30         r = NativeWasbRunner(NativeWasbConfig())
31         assert not isinstance(r, WslCommandMixin)
32         assert not hasattr(r, "_wsl_bash")
33
34
35     def test_reuses_model_agnostic_windowing():
36         from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
37         assert NativeWasbRunner.parse_trajectory_csv is TrackNetRunner.parse_trajectory_csv
38         assert NativeWasbRunner.trajectory_to_action_windows is
TrackNetRunner.trajectory_to_action_windows
39
40
41     # --- config resolution (env overrides win over config over default)
-----
42     def test_from_indexing_cfg_defaults():
43         cfg = NativeWasbConfig.from_indexing_cfg({})
44         assert cfg.device == "cuda" and cfg.python_bin == "python"
45         assert cfg.weights_path == "" and cfg.repo_dir == "~/models/WASB-SBDT"
46         assert cfg.stream_video is False and cfg.keep_frames is False
47
48
49     def test_from_indexing_cfg_reads_config_block():
50         cfg = NativeWasbConfig.from_indexing_cfg({"wasb": {
51             "weights_path": "/m/w.pth.tar", "device": "cpu", "python_bin":
"/envs/wasb/bin/python",
52             "repo_dir": "/m/WASB-SBDT", "stream_video": True, "keep_frames": True, "sport":
"tennis"}})
53         assert cfg.weights_path == "/m/w.pth.tar" and cfg.device == "cpu"
54         assert cfg.python_bin == "/envs/wasb/bin/python" and cfg.repo_dir == "/m/WASB-SBDT"
55         assert cfg.stream_video is True and cfg.keep_frames is True and cfg.sport ==
"tennis"
56
57
58     def test_env_overrides_beat_config(monkeypatch):
59         monkeypatch.setenv("WASB_WEIGHTS", "/env/w.pth.tar")
60         monkeypatch.setenv("WASB_DEVICE", "mps")
61         monkeypatch.setenv("WASB_PYTHON", "/env/py")
62         monkeypatch.setenv("WASB_REPO_DIR", "/env/repo")
63         monkeypatch.setenv("WASB_SCRATCH", "/env/scratch")
64         cfg = NativeWasbConfig.from_indexing_cfg({"wasb": {
65             "weights_path": "/cfg/w", "device": "cpu", "python_bin": "/cfg/py", "repo_dir":
"/cfg/repo"}})
66         assert cfg.weights_path == "/env/w.pth.tar" and cfg.device == "mps"
67         assert cfg.python_bin == "/env/py" and cfg.repo_dir == "/env/repo"
68         assert cfg.scratch_dir == "/env/scratch"
69
70
71     # --- healthcheck
-----
72     def _ok_cfg(tmp_path, **kw):
73         w = tmp_path / "w.pth.tar"
74         w.write_text("weights")
75         (tmp_path / "repo" / "src").mkdir(parents=True)
76         return NativeWasbConfig(weights_path=str(w), repo_dir=str(tmp_path / "repo"), **kw)
77
78
79     def test_healthcheck_ok(tmp_path):
80         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cuda"))
81         with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda
True")):

```

```

82         ok, msg = r.healthcheck()
83         assert ok and "torch" in msg
84
85
86     def test_healthcheck_empty_weights():
87         ok, msg = NativeWasbRunner(NativeWasbConfig(weights_path="")).healthcheck()
88         assert not ok and "weights" in msg.lower()
89
90
91     def test_healthcheck_weights_missing(tmp_path):
92         (tmp_path / "repo" / "src").mkdir(parents=True)
93         cfg = NativeWasbConfig(weights_path=str(tmp_path / "nope.pth"),
repo_dir=str(tmp_path / "repo"))
94         ok, msg = NativeWasbRunner(cfg).healthcheck()
95         assert not ok and "not found" in msg.lower()
96
97
98     def test_healthcheck_repo_missing(tmp_path):
99         w = tmp_path / "w.pth.tar"
100         w.write_text("x")
101         cfg = NativeWasbConfig(weights_path=str(w), repo_dir=str(tmp_path / "absent"))
102         ok, msg = NativeWasbRunner(cfg).healthcheck()
103         assert not ok and "repo" in msg.lower()
104
105
106     def test_healthcheck_torch_import_fails(tmp_path):
107         r = NativeWasbRunner(_ok_cfg(tmp_path))
108         with patch.object(r, "_run", return_value=_proc(1, stderr="ModuleNotFoundError:
torch")):
109             ok, msg = r.healthcheck()
110             assert not ok and "failed" in msg.lower()
111
112
113     def test_healthcheck_cuda_unavailable_when_device_cuda(tmp_path):
114         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cuda"))
115         with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda
False")):
116             ok, msg = r.healthcheck()
117             assert not ok and "cuda" in msg.lower()
118
119
120     def test_healthcheck_cpu_device_ok_without_cuda(tmp_path):
121         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cpu"))
122         with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda
False")):
123             ok, msg = r.healthcheck()
124             assert ok # cpu run does not require a GPU
125
126
127     def test_healthcheck_python_not_found(tmp_path):
128         r = NativeWasbRunner(_ok_cfg(tmp_path, python_bin="/no/such/python"))
129         with patch.object(r, "_run", side_effect=FileNotFoundError()):
130             ok, msg = r.healthcheck()
131             assert not ok and "python" in msg.lower()
132
133
134     # --- run_predict
-----
135     def _drive_success(cfg, tmp_path, vid="abc123abc123", **patches):
136         """Drive run_predict to success with the subprocess mocked; returns (runner, out,
rm_spy, argv, cwd)."""
137         r = NativeWasbRunner(cfg)
138         key = f"clip__{vid[:12]}"
139         (tmp_path / f"{key}_wasb.csv").write_text("Frame,Visibility,X,Y\n") # success = CSV
exists

```



```

140     with patch.object(r, "_copy_infer_script", return_value=True), \
141         patch.object(r, "_run", return_value=_proc(0)) as run_spy, \
142         patch.object(r, "_rm_rf") as rm_spy:
143         out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id=vid)
144         argv = run_spy.call_args[0][0]
145         cwd = run_spy.call_args.kwargs.get("cwd")
146         return r, out, rm_spy, argv, cwd, key
147
148
149     def test_run_predict_builds_native_argv_no_wsl(tmp_path):
150         _, out, rm, argv, cwd, key =
151         _drive_success(NativeWasbConfig(weights_path="/m/w.pth", device="cuda"), tmp_path)
152         assert out == os.path.join(os.path.abspath(str(tmp_path)), f"{key}_wasb.csv")
153         # No WSL / conda plumbing anywhere in the command.
154         assert "wsl" not in argv
155         assert not any("conda" in a or "/mnt/" in a for a in argv)
156         # It runs the SAME inference core, device-threaded, writing the SAME CSV name as
157         WSL.
158         assert argv[0] == "python" and argv[1] == "wasb_infer.py"
159         assert "--device" in argv and argv[argv.index("--device") + 1] == "cuda"
160         assert "--weights" in argv and argv[argv.index("--weights") + 1] == "/m/w.pth"
161         assert argv[argv.index("--out") + 1] == out
162         # Disk mode extracts frames + cleans them on success.
163         assert "--frames_out_dir" in argv and "--stream-video" not in argv
164         assert cwd.endswith(os.path.join("src"))
165         rm.assert_called_once()
166
167     def test_run_predict_stream_mode_no_frames_dir(tmp_path):
168         _, out, rm, argv, _cwd, _key = _drive_success(
169             NativeWasbConfig(weights_path="/m/w.pth", stream_video=True), tmp_path)
170         assert out is not None
171         assert "--stream-video" in argv and "--frames_out_dir" not in argv
172         rm.assert_not_called() # streaming never materializes a frame cache
173
174     def test_run_predict_device_threaded_cpu(tmp_path):
175         _, _out, _rm, argv, _cwd, _key = _drive_success(
176             NativeWasbConfig(weights_path="/m/w.pth", device="cpu"), tmp_path)
177         assert argv[argv.index("--device") + 1] == "cpu"
178
179
180     def test_run_predict_keep_frames_skips_cleanup(tmp_path):
181         _, _out, rm, _argv, _cwd, _key = _drive_success(
182             NativeWasbConfig(weights_path="/m/w.pth", keep_frames=True), tmp_path)
183         rm.assert_not_called()
184
185
186     def test_run_predict_none_when_copy_infer_fails(tmp_path):
187         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
188         with patch.object(r, "_copy_infer_script", return_value=False):
189             assert r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
190         is None
191
192     def test_run_predict_none_when_subprocess_fails(tmp_path):
193         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
194         with patch.object(r, "_copy_infer_script", return_value=True), \
195             patch.object(r, "_run", return_value=_proc(1, stderr="CUDA error")):
196             assert r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
197         is None
198
199     def test_run_predict_none_when_csv_absent(tmp_path):
200         # Subprocess "succeeds" (rc 0) but writes no CSV -> treated as failure.

```

```

201     r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
202     with patch.object(r, "_copy_infer_script", return_value=True), \
203         patch.object(r, "_run", return_value=_proc(0)):
204         assert r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
is None
205
206
207     def test_run_predict_times_out_returns_none(tmp_path):
208         import subprocess
209         r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth"))
210         with patch.object(r, "_copy_infer_script", return_value=True), \
211             patch.object(r, "_run", side_effect=subprocess.TimeoutExpired(cmd="x",
timeout=1)):
212             assert r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
is None
213
214
215     def test_csv_name_matches_wsl_contract(tmp_path):
216         """The native CSV name must equal the WSL runner's so downstream + the parity diff
line up."""
217         _, out, _rm, _argv, _cwd, key =
_drive_success(NativeWasbConfig(weights_path="/m/w.pth"), tmp_path)
218         assert os.path.basename(out) == f"{key}_wasb.csv"
219
220
221     # --- the detector_impl='native' seam (parallels the stub seam)
-----
222     def test_resolve_runner_native_selects_native_runner_for_wasb(monkeypatch):
223         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
224         from backend.pipeline.segmenters.wasb_hybrid import WasbHybridSegmenter
225         runner, params = WasbHybridSegmenter(None, {}). _resolve_runner({"detector_impl":
"native"})
226         assert isinstance(runner, NativeWasbRunner)
227         assert params.get("detector_impl") == "native"
228
229
230     def test_resolve_runner_native_env_overrides_config(monkeypatch):
231         monkeypatch.setenv("RALLY_DETECTOR_IMPL", "native")
232         from backend.pipeline.segmenters.wasb_hybrid import WasbHybridSegmenter
233         runner, _ = WasbHybridSegmenter(None, {}). _resolve_runner({"detector_impl": "auto"})
234         assert isinstance(runner, NativeWasbRunner)
235
236
237     def test_resolve_runner_native_supported_for_fusion_wasb(monkeypatch):
238         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
239         from backend.pipeline.segmenters.fusion_hybrid import FusionSegmenter
240         runner, params = FusionSegmenter(None, {}). _resolve_runner({"detector_impl":
"native"})
241         assert isinstance(runner, NativeWasbRunner)
242         assert params.get("fusion_substrate") == "wasb"
243
244
245     def test_resolve_runner_native_refused_for_fusion_tracknet(monkeypatch):
246         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
247         from backend.pipeline.segmenters.fusion_hybrid import FusionSegmenter
248         with pytest.raises(NotImplementedError):
249             FusionSegmenter(None, {}). _resolve_runner(
250                 {"detector_impl": "native", "fusion": {"substrate": "tracknet"}})
251
252
253     def test_resolve_runner_native_refused_for_tracknet(monkeypatch):
254         monkeypatch.delenv("RALLY_DETECTOR_IMPL", raising=False)
255         from backend.pipeline.segmenters.tracknet_hybrid import TrackNetHybridSegmenter
256         with pytest.raises(NotImplementedError):
257             TrackNetHybridSegmenter(None, {}). _resolve_runner({"detector_impl": "native"})

```

6. assistant (2026-06-20T09:00:05.705Z)

Now let me verify the claim. The finding says the tests assert `--frames\_out\_dir` is present but never assert its VALUE, and never set scratch_dir. Let me check the test assertions carefully and compare against the actual code path. Let me also check the existing WSL runner for the equivalent contract and see how the WSL runner is tested for the scratch path.

7. user (2026-06-20T09:00:07.125Z)

```

1      """WASB shuttle-trajectory runner (WSL2 subprocess adapter).
2
3      windows-side counterpart to backend/detectors/wasb_infer.py. It:
4          1. stages the video into the WSL filesystem (fast local reads, not /mnt),
5          2. copies wasb_infer.py into the WASB repo's src/ dir,
6          3. runs it in the WSL `wasb` conda env -> a Frame,Visibility,X,Y trajectory CSV,
7          4. reuses the (model-agnostic) trajectory->action-window logic from tracknet_runner.
8
9      So the only WASB-specific bits live here; the windowing + CSV parsing are shared
10     with the TrackNet path. Verified inference core: wasb_infer.py (see its docstring).
11     """
12     import os
13     import shlex
14     import logging
15     import subprocess
16     from dataclasses import dataclass
17     from typing import Dict, Any, Optional, Tuple
18
19     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical.
20     # parse_trajectory_csv / trajectory_to_action_windows are @staticmethods on
21     TrackNetRunner.
22     from backend.pipeline.detectors.tracknet_runner import to_wsl_mnt_path, TrackNetRunner
23     from backend.pipeline.detectors.base import DetectorRunner, WslCommandMixin,
24     wsl_tilde_quote
25
26     logger = logging.getLogger(__name__)
27
28     # Windows path to the inference wrapper we stage into WSL.
29     _INFER_WIN = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
30
31     @dataclass
32     class WasbConfig:
33         """Resolved settings for a WASB WSL run (built from config.json ->
34         indexing.wasb)."""
35         distro: str = "Ubuntu"
36         repo_dir: str = "~/models/WASB-SBDT"
37         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
38         conda_env: str = "wasb"
39         weights_path: str = "~/models/WASB-
40         SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
41         sport: str = "badminton"
42         wsl_stage_dir: str = "~/clips"
43         timeout_sec: int = 1800 # 30 min; a hung GPU/conda call is killed (0 = no timeout)
44         keep_frames: bool = False # retain staged video + frame cache after success (debug
45         only)
46         min_free_gb: float = 0.0 # warn if WSL free space below this before a run (0 =
47         off)
48         # FAST PATH (default OFF until verified): decode the staged video on the fly inside
49         # WSL and feed frames straight to the detector, skipping the slow per-frame PNG
50         # extraction that dominates a long clip. Output is bit-identical to the PNG path
51         # (PNG is lossless), but this stays opt-in until the owner's side-by-side confirms
52         it.
53         stream_video: bool = False

```

```

48
49 @classmethod
50 def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "WasbConfig":
51     w = (idx_cfg or {}).get("wasb", {}) or {}
52     return cls(
53         distro=w.get("wsl_distro", "Ubuntu"),
54         repo_dir=w.get("repo_dir", "~/models/WASB-SBDT"),
55         conda_sh=w.get("conda_sh", "~/miniconda3/etc/profile.d/conda.sh"),
56         conda_env=w.get("conda_env", "wasb"),
57         weights_path=w.get("weights_path",
58                             "~/models/WASB-
SBDT/pretrained_weights/wasb_badminton_best.pth.tar"),
59         sport=w.get("sport", "badminton"),
60         wsl_stage_dir=w.get("wsl_stage_dir", "~/clips"),
61         timeout_sec=int(w.get("timeout_sec", 1800)),
62         keep_frames=bool(w.get("keep_frames", False)),
63         min_free_gb=float(w.get("min_free_gb", 0.0)),
64         stream_video=bool(w.get("stream_video", False)),
65     )
66
67
68 class WasbRunner(WslCommandMixin, DetectorRunner):
69     def __init__(self, cfg: WasbConfig):
70         self.cfg = cfg
71
72     def healthcheck(self) -> Tuple[bool, str]:
73         """Verify the WSL `wasb` env imports torch and sees a GPU. Returns (ok, msg)."""
74         cmd = (f"conda activate {self.cfg.conda_env} && "
75              f"python -c \"import torch; print('torch', torch.__version__, "
76              f"'cuda', torch.cuda.is_available())\"")
77         try:
78             res = self._wsl_bash(cmd)
79         except FileNotFoundError:
80             return False, "`wsl` not found ? Windows + WSL2 required."
81         except subprocess.TimeoutExpired:
82             return False, "WSL healthcheck timed out."
83         out = (res.stdout or "").strip()
84         if res.returncode != 0:
85             return False, f"WSL `wasb` env check failed: {(res.stderr or out).strip()}"
86         return True, out
87
88     def _stage_video(self, video_win_path: str, dest_name: str) -> Optional[str]:
89         src_mnt = to_wsl_mnt_path(video_win_path)
90         reason = self.wsl_source_unreadable_reason(src_mnt)
91         if reason:
92             logger.error("Cannot stage video into WSL: %s", reason)
93             return None
94         dest = f"{self.cfg.wsl_stage_dir}/{dest_name}"
95         cmd = f"mkdir -p {self.cfg.wsl_stage_dir} && cp {shlex.quote(src_mnt)}
{wsl_tilde_quote(dest)} && echo OK"
96         try:
97             res = self._wsl_bash(cmd)
98         except subprocess.TimeoutExpired:
99             logger.error("Staging video into WSL timed out.")
100             return None
101         if res.returncode != 0 or "OK" not in (res.stdout or ""):
102             logger.error("Failed to stage video into WSL: %s", (res.stderr or
res.stdout).strip())
103             return None
104         return dest
105
106     def _stage_infer_script(self) -> bool:
107         """Copy wasb_infer.py into the WASB repo src/ so it can import WASB modules."""
108         src_mnt = to_wsl_mnt_path(_INFER_WIN)
109         cmd = f"cp {shlex.quote(src_mnt)} {self.cfg.repo_dir}/src/wasb_infer.py && echo

```

```

OK"
110         res = self._wsl_bash(cmd)
111         return res.returncode == 0 and "OK" in (res.stdout or "")
112
113     def run_predict(self, video_win_path: str, output_win_dir: str,
114                    video_id: Optional[str] = None) -> Optional[str]:
115         """Run WASB on a video. Returns the Windows path to the trajectory CSV, or None.
116
117         All WSL/cache/output artifacts are namespaced by ``video_id`` (the file MD5) so
118         two videos that share a filename (e.g. two ``match.mp4`` from different folders)
119         can never reuse each other's staged frames, resumable cache, or CSV.
120
121         # Resolve relative dirs (the hybrids pass "output/wasb") BEFORE the /mnt
122         # translation: to_wsl_mnt_path passes relative paths through unchanged, so
123         # WSL resolved them against the inference cwd (~/.models/WASB-SBDT/src) and
124         # the CSV landed inside WSL while Windows polled the repo-relative path.
125         output_win_dir = os.path.abspath(output_win_dir)
126         os.makedirs(output_win_dir, exist_ok=True)
127         # Normalize for cross-platform basename (tests use Windows paths on Linux).
128         video_win_path = str(video_win_path).replace("\\", "/")
129         base = os.path.basename(video_win_path)
130         stem, ext = os.path.splitext(base)
131         # Unique, content-derived key: same filename + different bytes -> different key.
132         key = f"{stem}__{video_id[:12]}" if video_id else stem
133         wsl_out_dir = to_wsl_mnt_path(output_win_dir)
134         expected_csv_win = os.path.join(output_win_dir, f"{key}_wasb.csv")
135
136         if not self._stage_infer_script():
137             logger.error("Could not stage wasb_infer.py into WSL.")
138             return None
139         staged_video = self._stage_video(video_win_path, f"{key}{ext}")
140         if staged_video is None:
141             return None
142         frames_dir = f"{self.cfg.wsl_stage_dir}/{key}_frames"
143
144         # Disk guard: a 4K clip can extract ~100 GB of PNG frames. Warn early if the
145         # WSL stage filesystem is already low so the user can clean up before it fills.
146         # (Streaming never materializes the PNGs, so the guard only matters off the fast
147         path.)
148         if self.cfg.min_free_gb > 0 and not self.cfg.stream_video:
149             free = self._wsl_free_gb(self.cfg.wsl_stage_dir)
150             if free is not None and free < self.cfg.min_free_gb:
151                 logger.warning(
152                     "Low WSL disk: %.1f GB free at %s (< min_free_gb=%.1f). Frame
153                     extraction "
154                     "may fill the disk ? consider running tools/cleanup_caches.py.",
155                     free, self.cfg.wsl_stage_dir, self.cfg.min_free_gb)
156
157         # Fast path: stream-decode the video in-process (no PNG extraction). Output is
158         # bit-identical to the PNG round-trip but avoids writing ~46k PNGs for a 12-min
159         clip.
160         if self.cfg.stream_video:
161             frame_args = "--stream-video "
162         else:
163             frame_args = f"--frames_out_dir {wsl_tilde_quote(frames_dir)} "
164         cmd = (
165             f"conda activate {self.cfg.conda_env} && cd {self.cfg.repo_dir}/src && "
166             f"python wasb_infer.py "
167             f"--video {wsl_tilde_quote(staged_video)} "
168             f"{frame_args}"
169             f"--weights {wsl_tilde_quote(self.cfg.weights_path)} "
170             f"--sport {shlex.quote(self.cfg.sport)} "
171             f"--out {wsl_tilde_quote(wsl_out_dir + '/' + key + '_wasb.csv')}"
172         )
173         logger.info("Running WASB inference on %s ...", base)

```

```

171         try:
172             res = self._wsl_bash(cmd)
173         except subprocess.TimeoutExpired:
174             logger.error("WASB inference timed out after %ss.", self.cfg.timeout_sec)
175             return None
176         if res.returncode != 0:
177             logger.error("WASB inference failed:\n%s", (res.stderr or
res.stdout)[-2000:])
178             return None
179         if not os.path.exists(expected_csv_win):
180             logger.error("WASB finished but expected CSV not found at %s",
expected_csv_win)
181             return None
182         logger.info("WASB trajectory CSV: %s", expected_csv_win)
183
184         # Success ? the trajectory CSV is the durable output; the staged video copy and
185         # the (huge) decoded frame cache are now pure intermediates. Delete them unless
186         # explicitly retained. On any earlier failure we return above WITHOUT cleaning,
187         # so the frame cache + detector cache remain for a resumable re-run.
188         # (Streaming never writes a frame cache, so there's only the staged video to
drop.)
189         if not self.cfg.keep_frames:
190             stale = [staged_video] if self.cfg.stream_video else [frames_dir,
staged_video]
191             logger.info("Cache hygiene: removing staged video%s for %s "
"(set indexing.wasb.keep_frames=true to retain).",
192                         "" if self.cfg.stream_video else " + frame cache", key)
193             self._wsl_rm_rf(stale)
194         return expected_csv_win
195
196
197         # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
198         parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
199         trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)
200

```

8. user (2026-06-20T09:00:07.917Z)

```

test_calibrate_wasb.py
test_config_wasb.py
test_export_wasb_segments.py
test_native_wasb_runner.py
test_wasb_cache.py
test_wasb_runner.py

```

9. assistant (2026-06-20T09:00:11.825Z)

Now let me verify the precise claims in the finding by examining the test assertions closely. The finding makes specific claims:

1. Tests assert ``--frames_out_dir` is PRESENT but never assert its VALUE.
2. Tests never set `scratch\_dir`, so the env-driven scratch path is never followed end-to-end into argv.
3. The frames dir that gets `\_rm\_rf`'d on cleanup is never asserted (the `\_rm\_spy` arg is never checked).

Let me grep the test file for these assertions to confirm.

10. user (2026-06-20T09:00:18.201Z)

```

63: monkeypatch.setenv("WASB_SCRATCH", "/env/scratch")
68: assert cfg.scratch_dir == "/env/scratch"
136: """Drive run_predict to success with the subprocess mocked; returns (runner, out,

```

```

rm_spy, argv, cwd)."""
142:         patch.object(r, "_rm_rf") as rm_spy:
144:     argv = run_spy.call_args[0][0]
145:     cwd = run_spy.call_args.kwargs.get("cwd")
146:     return r, out, rm_spy, argv, cwd, key
161:     assert "--frames_out_dir" in argv and "--stream-video" not in argv
163:     rm.assert_called_once()
170:     assert "--stream-video" in argv and "--frames_out_dir" not in argv
171:     rm.assert_not_called() # streaming never materializes a frame cache
183:     rm.assert_not_called()

```

11. assistant (2026-06-20T09:00:26.791Z)

The grep confirms the finding's claims about the test gaps. Let me verify each precise claim:

```

**Claim 1: Tests assert `--frames_out_dir` is PRESENT but never assert its VALUE.**
- Line 161: `assert "--frames_out_dir" in argv and "--stream-video" not in argv` ? presence
only.
- Line 170: `assert "--stream-video" in argv and "--frames_out_dir" not in argv` ? absence
(stream mode).
- No test does `argv[argv.index("--frames_out_dir") + 1] == <value>`. Confirmed.

**Claim 2: Tests never set `scratch_dir` in a run_predict drive.**
- Line 68: `assert cfg.scratch_dir == "/env/scratch"` ? this only tests config *resolution*
(`from_indexing_cfg`), NOT that the resolved scratch_dir flows into the argv `--frames_out_dir`
value via run_predict.
- The `_drive_success` helper (lines 135-146) never passes `scratch_dir`. All run_predict tests
use `NativeWasbConfig(weights_path=..., ...)` with `scratch_dir` defaulting to `""`. Confirmed.

**Claim 3: The `_rm_rf` target value is never asserted.**
- Line 163: `rm.assert_called_once()` ? call count only.
- Line 171/183: `rm.assert_not_called()`.
- No test checks `rm.call_args` to confirm the deleted path equals the frames_dir passed as
`--frames_out_dir`. Confirmed.

```

Now let me verify the WHY ? the failure scenario. Let me check the actual code path at lines 169-171 and 198-201 to confirm the scratch_dir ? frames_dir ? argv ? cleanup wiring is real and that a regression would be undetected.

12. user (2026-06-20T09:00:27.581Z)

```

164         frames_dir: Optional[str] = None
165         if self.cfg.stream_video:
166             # Fast path: decode in-process, no PNG extraction. Output is bit-identical
to disk.
167             argv.append("--stream-video")
168         else:
169             scratch = self.cfg.scratch_dir or output_dir
170             frames_dir = os.path.join(scratch, f"{key}_frames")
171             argv += ["--frames_out_dir", frames_dir]
172             argv += ["--weights", weights,
173                     "--sport", self.cfg.sport,
174                     "--device", self.cfg.device,
175                     "--out", out_csv]
176
177             logger.info("Running native WASB inference (device=%s) on %s ...",
178                         self.cfg.device, os.path.basename(norm))
179             try:
180                 res = self._run(argv, cwd=self._repo_src())
181             except FileNotFoundError:
182                 logger.error("python binary not found: %r (set WASB_PYTHON).",
self.cfg.python_bin)
183                 return None
184             except subprocess.TimeoutExpired:
185                 logger.error("WASB native inference timed out after %ss.",

```

```

self.cfg.timeout_sec)
186         return None
187         if res.returncode != 0:
188             logger.error("WASB native inference failed:\n%s", (res.stderr or res.stdout
or "")[-2000:])
189         return None
190         if not os.path.exists(out_csv):
191             logger.error("WASB native finished but expected CSV not found at %s",
out_csv)
192         return None
193         logger.info("WASB trajectory CSV: %s", out_csv)
194
195         # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
196         # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
197         # (Streaming never materializes frames, so there is nothing to drop on that
path.)
198         if not self.cfg.keep_frames and frames_dir:
199             logger.info("Cache hygiene: removing frame cache for %s "
200                         "(set indexing.wasb.keep_frames=true to retain).", key)
201             self._rm_rf(frames_dir)
202         return out_csv

```

13. assistant (2026-06-20T09:00:32.198Z)

The code is exactly as the finding describes. Now let me assess the **severity** and the "real failure scenario" rigorously, since the finding is a **test-quality** (coverage gap) finding, not a code bug. Let me confirm the regression it describes would truly stay green by mentally tracing a hypothetical regression, and verify there's no OTHER test that would catch it. Let me run the actual test file to confirm current state, and check whether any test indirectly covers the scratch path value.

14. user (2026-06-20T09:00:41.161Z)

```

..... [100%]
28 passed in 6.01s

```

15. assistant (2026-06-20T09:00:48.194Z)

All 28 tests pass. Now let me confirm the regression scenarios the finding describes would actually slip through. Let me trace both hypothetical regressions:

Regression A: ``scratch = self.cfg.scratch_dir or output_dir`` ? drops scratch, becomes just ``output_dir``. **In that case**, with the default config (``scratch_dir=""``), ``scratch`` would already be ``output_dir`` anyway ? so default-config tests are blind by construction. With a configured ``scratch_dir='/scratch'``, the frames_dir would wrongly land under output_dir. But **no run_predict test sets scratch_dir**, so this regression is undetected. Confirmed.

Regression B: cleanup target diverges from extraction target. **If** ``_rm_rf(frames_dir)`` were changed to delete a different/wrong path, line 163's ``rm.assert_called_once()`` would still pass (it only checks call count, not the argument). Confirmed.

Let me double-check there's no value assertion I missed by reading the full assertion block of the relevant test once more, and verify the WSL runner's test suite for comparison (to confirm this is a genuine coverage delta, not the repo's accepted norm).

16. user (2026-06-20T09:00:50.159Z)

```

57:# --- cache hygiene (keep_frames / min_free_gb / shared WSL helpers) -----
60:     cfg = WasbConfig.from_indexing_cfg({"wasb": {"keep_frames": True, "min_free_gb": 25}})
61:     assert cfg.keep_frames is True and cfg.min_free_gb == 25.0
63:     assert default.keep_frames is False and default.min_free_gb == 0.0
66: def test_wsl_rm_rf_builds_rm_command():

```



```

69:         r._wsl_rm_rf(["~/clips/a_frames", "~/clips/a.mp4"])
70:     inner = wb.call_args[0][0]
71:     assert inner.startswith("rm -rf ") and "a_frames" in inner and "a.mp4" in inner
74: def test_wsl_rm_rf_noop_on_empty():
77:     r._wsl_rm_rf([])
97:     patch.object(r, "_wsl_rm_rf") as rm:
103: def test_run_predict_cleans_frames_and_staged_video_on_success(tmp_path):
105:     rm.assert_called_once()
106:     purged = rm.call_args[0][0]
107:     assert f"~/clips/{key}_frames" in purged and f"~/clips/{key}.mp4" in purged
110: def test_run_predict_keeps_frames_when_keep_frames_true(tmp_path):
111:     rm, _ = _run_predict_to_success(WasbConfig(keep_frames=True), tmp_path)

```

17. assistant (2026-06-20T09:00:56.080Z)

This is the key comparison. The **WSL runner's test suite** (`test_wasb_runner.py`) ? the established sibling this native runner is modeled on ? DOES assert the cleanup target value:

- Line 105: `rm.assert_called_once()`
- Lines 106-107: `purged = rm.call_args[0][0]` then `assert f"~/clips/{key}_frames" in purged` and `f"~/clips/{key}.mp4" in purged`

So the repo's own convention/precedent for the sibling runner **does** assert the `_rm_rf` argument value, while the new native runner test only asserts `rm.assert_called_once()` (line 163) and never inspects `rm.call_args`. This confirms the finding is identifying a genuine coverage regression relative to the established sibling pattern.

Let me also confirm the WSL test asserts the frames path in the command argv (the equivalent of the `--frames_out_dir` value check) to gauge whether the value-assertion gap is also a delta.

18. user (2026-06-20T09:00:57.065Z)

```

80
81     def test_wsl_free_gb_pares_bytes_else_none():
82         r = WasbRunner(WasbConfig())
83         with patch.object(r, "_wsl_bash", return_value=_proc(0, stdout=str(50 * 1024 ** 3) +
84 "\n")):
85             assert abs(r._wsl_free_gb("~/clips") - 50.0) < 0.01
86             with patch.object(r, "_wsl_bash", return_value=_proc(0, stdout="garbage")):
87                 assert r._wsl_free_gb("~/clips") is None
88
89     def _run_predict_to_success(cfg, tmp_path, vid="abc123abc123"):
90         """Drive run_predict to its success branch with WSL fully mocked. Returns (rm_spy,
91 key)."""
92         r = WasbRunner(cfg)
93         key = f"clip_{vid[:12]}"
94         (tmp_path / f"{key}_wasb.csv").write_text("Frame,Visibility,X,Y\n") # success = CSV
95         exists
96         with patch.object(r, "_stage_infer_script", return_value=True), \
97             patch.object(r, "_stage_video", return_value=f"~/clips/{key}.mp4"), \
98             patch.object(r, "_wsl_bash", return_value=_proc(0)), \
99             patch.object(r, "_wsl_rm_rf") as rm:
100             out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id=vid)
101             assert out is not None
102             return rm, key
103
104     def test_run_predict_cleans_frames_and_staged_video_on_success(tmp_path):
105         rm, key = _run_predict_to_success(WasbConfig(), tmp_path)
106         rm.assert_called_once()
107         purged = rm.call_args[0][0]
108         assert f"~/clips/{key}_frames" in purged and f"~/clips/{key}.mp4" in purged
109

```

```

110     def test_run_predict_keeps_frames_when_keep_frames_true(tmp_path):
111         rm, _ = _run_predict_to_success(WasbConfig(keep_frames=True), tmp_path)
112         rm.assert_not_called()
113

```

19. assistant (2026-06-20T09:01:12.194Z)

I now have everything I need. My verification is complete.

Summary of verification

****The finding's factual claims are all confirmed against the working-tree code:****

1. ****`--frames\_out\_dir` value is never asserted.**** `tests/test\_native\_wasb\_runner.py:161` asserts only `"--frames_out_dir"` in `argv``. No test does `argv[argv.index("--frames_out_dir")+1] == <expected>``. Confirmed.

2. ****`scratch\_dir` is never exercised through `run\_predict`.** The ``_drive_success`` helper (lines 135-146) never passes ``scratch_dir``; every `run_predict` test leaves it at default `""``. The only scratch coverage (line 68) tests ``from_indexing_cfg`` resolution, not the ``scratch ? frames_dir ? argv`` path in ``run_predict`` (`native_wasb_runner.py:169-171`). With the default ``scratch_dir=""``, ``scratch = self.cfg.scratch_dir` or `output_dir`` already collapses to ``output_dir``, so a regression that drops ``scratch_dir`` from the join is invisible to the existing tests. Confirmed.

3. ****The ``rm_rf`` cleanup target value is never asserted.**** Line 163 is ``rm.assert_called_once()`` (count only); ``rm.call_args`` is never inspected. A cleanup target that diverged from the extraction ``frames_dir`` would leave the test green. Confirmed.

****Failure scenario is real and concrete:**** `native_wasb_runner.py` targets long 4K clips on a Linux GPU box; ``scratch_dir`` is a documented first-class env contract (``WASB_SCRATCH`/`RALLY_SCRATCH`/`TMPDIR``, module docstring lines 15-17, config line 77) whose entire purpose is to keep the GBs of extracted PNG frames off the small system disk. A regression dropping ``scratch`` from the path join, or pointing ``rm_rf`` at the wrong dir, is a disk-fill / leaked-frames regression that the current tests cannot catch.

****Repo precedent corroborates this is a genuine coverage gap, not the accepted norm:**** the sibling WSL runner ? which this native runner is explicitly modeled on ? DOES assert the cleanup target value (``test_wasb_runner.py:106-107``: ``purged = rm.call_args[0][0]; assert f"~/clips/{key}_frames" in purged ...``). The new native test regressed away that assertion.

****Severity calibration:**** This is a test-coverage gap, not a live code bug ? the production code at lines 169-171 and 198-201 is currently correct, and the byte-parity acceptance gate is explicitly owner-verified on the GPU box (out of scope for this suite). The undetected-regression risk is real and matches a documented disk-fill failure mode the repo cares about, but it is latent (requires a future edit to manifest) and the proposed fix is a small additive test. That fits "minor" better than "major" ? the reviewer's "major" overstates it given the code is correct today and the fix is purely additive test hardening.

20. user (2026-06-20T09:01:21.710Z)

Structured output provided successfully

21. assistant (2026-06-20T09:01:27.843Z)

Verification complete. The finding is ****real**** but I've corrected the severity from major to ****minor****.

All three of the finding's factual claims hold against the working tree:

- ``tests/test_native_wasb_runner.py:161`` asserts only the presence of ``--frames_out_dir``, never its value.
- No ``run_predict`` test sets ``scratch_dir``; the only scratch coverage (line 68) tests config resolution, not the ``scratch ? frames_dir ? argv`` wiring at ``native_wasb_runner.py:169-171``. With the default ``scratch_dir=""``, ``scratch` or `output_dir`` already collapses to ``output_dir``, so

a regression dropping ``scratch_dir`` from the join stays green.

- Cleanup is asserted by count only (line 163 ``rm.assert_called_once()``); ``rm.call_args`` is never inspected.

The repo's own precedent corroborates this is a genuine gap: the sibling WSL runner's test (``tests/test_wasb_runner.py:106-107``) DOES assert the ``_rm_rf`` target value, which the new native test regressed away.

I downgraded to minor because the production code is correct today, the byte-parity gate is owner-verified on the GPU box (out of scope for this suite), and the risk is latent ? it requires a future edit to manifest, and the fix is a small additive test.