

claude-windows / 9f6357d8-0e35-491a-a830-052e094c7f37

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\9f6357d8-0e35-491a-a830-052e094c7f37\subagents\workflows\wf\_b6ef6275-de5\agent-a711967821136d6a9.jsonl`
- SHA-256: `39feaa8c830ec6d7f6ea159eeb944832ebb7c72050db0a08c457447ddc3ac3b3`
- Source modified: `2026-06-21T13:58:45+00:00`
- Imported at: `2026-07-05T16:48:25+00:00`
- project: `wf\_b6ef6275-de5`
- session_id: `9f6357d8-0e35-491a-a830-052e094c7f37`

Transcript

1. user (2026-06-21T13:56:13.270Z)

Read the RALLY-DETECTION QUALITY track docs in the engine repo (C:/Users/avidu/Projects/badminton-highlight-indexer) and extract status + next steps. Read docs/RALLY_QUALITY_RESEARCH.md (esp. its §7 progress tracker), docs/HOW_RALLY_DETECTION_WORKS.md, docs/MULTI_SIGNAL_FUSION_PLAN.md (the held-shuttle Gate-A rally_gate.py not wired into production = cheapest win), docs/EXPERIMENT_HARNESS.md, docs/GOLDEN_REGRESSION_FIXTURES.md, docs/ROORA_LABELING_GUIDELINES.md, and docs/NEXT_STEPS.md (the prioritized pick-up for this track). Capture the prioritized next steps (R2 metric ? R4 end-rule, Gate-A wiring, golden-set growth, etc.), gates, effort. Return structured JSON (track="rally-quality").

2. assistant (2026-06-21T13:56:16.232Z)

I'll start by pulling the repo to ensure freshness, then read the rally-quality track docs.

3. user (2026-06-21T13:56:19.661Z)

fatal: Cannot fast-forward to multiple branches.

4. user (2026-06-21T13:56:19.880Z)

ABLATION_LAYER.md
ANALYZERS_AND_RECONCILERS.md
BACKLOG.md
CLOUD_GPU_DRYRUN_GUIDE.md
CODE_AUDIT_AND_TEST_HARDENING.md
CODE_MAP.md
COMMERCIALIZATION.md
COMPUTE_DECOUPLED_SERVING
DEFERRED_HARDENING_PLAN.md
DESKTOP_APP_PLAN.md
DOC_STATUS.md
EVALUATION.md
EXECUTION_POLICY.md
EXPERIMENT_HARNESS.md
GOLDEN_REGRESSION_FIXTURES.md
HARDENING_LOOP_HANDOFF.md
HOSTING_PLAN.md
HOW_RALLY_DETECTION_WORKS.md
I18N_PLAN.md
KHELSUTRA_PRODUCT_OVERVIEW.md
MULTI_SIGNAL_FUSION_PLAN.md
NEXT_STEPS.md
OWNED_MODEL_IMPLEMENTATION_PLAN.md
PERSONALIZATION_PLAN.md

PER_VIDEO_WORKSPACE.md
PLATFORM_ARCHITECTURE.md
POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md
PROJECT_DOC_TEMPLATE.md
PROMINENT_COURT_DETECTION.md
QUALITY_ITERATIONS.md
R1_MILESTONE_LAUNCH_REPORT.md
R2_EVAL_METRIC_DESIGN.md
RALLY_DETECTION_GAP_CLOSURE.md
RALLY_DETECTION_QUALITY_REPORT.md
RALLY_QUALITY_RESEARCH.md
README.md
REAL_FOOTAGE_VALIDATION_REPORT.md
ROADMAP.md
SETUP_NEW_MACHINE.md
STORAGE_SHARING_MODEL.md
TRACKNET_WSL_SETUP.md
UI_ALPHA_EXECUTION_PLAN.md
UI_PROPOSAL.md
VIDEO_LOCALITY_MODEL.md
archives
data_pipeline
mockups

5. assistant (2026-06-21T13:56:22.476Z)

The branch has multiple upstreams; I'll proceed read-only since I cannot cleanly fast-forward. Let me read the rally-quality track docs.

6. user (2026-06-21T13:56:23.473Z)

```
1      # Next steps (live) ? refreshed 2026-06-21
2
3      > **How to read:** this top section is the **current cross-track state + prioritized
pending list** (refreshed 2026-06-21, replacing the stale 2026-06-13 header). The detailed
**owned-model & rally-quality track** plan (2026-06-13) is **preserved below, still valid** ? it
is now *one of several* active tracks. Live blow-by-blow: memory `current-state.md`. Doc
lifecycle/hygiene: `docs/DOC_STATUS.md`.
4
5      ## ? Where we are (2026-06-21)
6      **DECOUPLED_COMPUTE delivered on GCP + archived** (#270); **compute = GCP only**
(ADR-013, DO dropped). **`COMPUTE_DECOUPLED_SERVING`** subproject spun up (#272, **ADR-014**) ?
one pure core + local/cloud profiles; **owner M0 design sign-off pending**. **Per-video
workspace P0** shipped (#264); **doc-status register + archival due-diligence** in place (#273).
The **khelsutra site + per-rally UI** ship in parallel (sibling track). The **rally-quality /
owned-model** track (below) is the core ML work, gated on **growing the golden corpus**.
7
8      ## ? Prioritized pending ? all tracks
9      **P0 ? highest leverage / unblocks the most**
10     1. ? **`COMPUTE_DECOUPLED_SERVING` M0 approved + M1?M4 SHIPPED (#279/#280/#282/#283) +
M5+ batch authorized** (owner 2026-06-21, ADR-014 ratified). All $8 gating Qs locked (D7?D17).
**Owner residual** (not blocking Claude's default-OFF code): counsel ToS/DPA (M9-consent live +
M11 live), real GCP prices+FX (M8 calibrate), Google OAuth app (M12 real), CF team-domain/AUD +
ingress-lock; real cloud verification authorized within a **$100/?10k cap** (D17).
11     2. **[owner, ongoing] Grow the golden corpus** ? multi-court badminton + ?3
matches/sport (TT first). The single highest-leverage move; the quality track *and* full-corpus
training feed on it.
12     3. **[owner + Claude · GPU] Full-corpus training (n?5)** ? a real `weights/<ver>` (the
`0.1.0-l4` bundle is n=2 plumbing-only).
13
14     **P1 ? near-term, mostly Claude-doable (default-OFF, green-gated)**
15     4. **[Claude] `COMPUTE_DECOUPLED_SERVING` M5?M12 default-OFF code batch** (owner-
authorized 2026-06-21, merge-each-on-green): M5 truly-local preset+startup-checks+L4 parity test
? M6 `backend/compute/` Cloud Run dispatch + Dockerfile (mocked) ? M8 v6 ledger+pricebook+/cost
? M9-auth Cf-Access JWT+tenant-scoping ? M11 flywheel plumbing ? M12 `gdrive-api` backend.
```

```

*(M1?M4 config/input/output/device layer = ? SHIPPED.)*
16      5. **[Claude] Per-video workspace P1?P6** (P1 v6 DB migration ? P2 worker ? P3 per-video
routing of `detector/`+`tmp/` *[the `output/chunks_` de-hardcode base = ? done via #282]* ? P4
? P5 Launch-Report flip ? P6 GPU parity). ? P1's v6 migration coordinates with the M8 v6 cost-
ledger migration ? author the version bump once.
17      6. **[Claude] Ops-Agent codification into `deploy/gcp`** (every VM auto-enables
observability ? [[gcp-ops-agent-enable]]).
18      7. **[owner] Set the ship-gate target (#172)** ? needs corpus growth.
19      8. **[owner-side GPU] R-series: promote-or-reject `R2` tolF1 as the gate** (the
ablation) ? needs GPU-extracted golden features.
20
21      **P2 ? later / gated / deferred**
22      9. Rally-quality **P3/P4** (person-fusion LOVO, audio cue) ? needs GPU golden features.
23      10. **M0.4 scorer swap** ? TemporalMaxer (per the 2026-06-17 external review).
24      11. **Multisport** (table-tennis first, then pickleball).
25      12. Serving **M8/M9** ? per-video cost ledger, edge auth, DPA/consent (within
`COMPUTE_DECOUPLED_SERVING`; owner/legal-gated).
26      13. **WASB-C3** weight provenance (gates *sharing* a trained model) ? legal.
27      14. **Rename repo off "badminton" (owner, BACKLOG).
28      15. Owner housekeeping: retire the shared **D0 droplet** + **D0 token**; collector
**`md5`** export (collector repo).
29      16. **PMF field validation** (owner-side, business track).
30
31      **Doc-hygiene (tracked in `docs/DOC_STATUS.md` $4):** ? #1 per-video reconcile + ? #2
this refresh; **pending:** field/PMF/Roora ? archives + a data-layer `docs/` dir (owner boundary
confirm); confirm the COMPLETED top-level stubs; verify `HOSTING_PLAN`/`UI_PROPOSAL` currency.
32
33      ---
34
35      ## Owned-model & rally-quality track ? detail (2026-06-13 plan; still valid)
36
37      **Updated:** 2026-06-13 (**rally-detection quality track**: fusion + experiment-harness
design docs merged,
38      owner-gated ? see the dedicated section below, the desktop-agent pick-up). Prior:
2026-06-10 (audit?M0?Gemini-battery session; later same day the **UI-alpha track opened**:
39      ADR-009 ratified ? KhelSutra Guru / khelsutra.guru, private `khelsutra` repo, local-
first delivery ? see
40      `docs/UI_ALPHA_EXECUTION_PLAN.md`; **PR-1 async jobs (#16) MERGED 2026-06-11, PR #87**;  

**PR-2
41      upload/download MERGED 2026-06-11, PR #99** ? live-E2E-verified (3-part chunked upload
MD5 byte-identity +
42      compile?browser-download); next: PR-3 identity. Same day: the **8?9 quality sweep
#90?#98 landed** ? ruff/mypy/
43      coverage-floor CI lanes, SQLite conn-close fix, config unknown-key warning, hybrid-twin
collapse into
44      `trajectory_hybrid.py`, WSL tilde+abspath live-path fixes; suite 388 green; LOVO 0.626
re-verified exact).
45      **Authoritative roadmap:**
46      `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` (+ ADR-008 for topology). **Latest blow-by-
blow:** memory
47      `current-state.md` / `session-handoff.md`. This file = the prioritized "what to do
next," refreshed each session.
48
49      ## ? Next-session / GPU-box pickup (2026-06-15)
50      **M0 in-container quick wins are COMPLETE:** ship-gate provisional blocker (**#170**),
train/eval split guard (**#171**), the **doc-consistency sweep** (**#174** ? corrected 16
drifted docs vs the code), and the **P3 person-feature fusion litmus machinery is already
present + suite-green** (`backend/eval/fusion_compare.py` + `ablation.py`: person-driven ?F1 /
bootstrap CI / paired sign test; the real-golden `skipif` litmus waits only on GPU-extracted
data). ? CI is red repo-wide due to a **GitHub Actions minutes/billing outage** (account-level,
not code) ? work is local-gated + `--admin`-merged until Actions is restored (Settings ? Billing
? Actions minutes / spending limit).
51
52      **Immediate next step ? the real P3 LOVO measurement** (this box has a GPU +
`torch+cu124` + `cv2`; videos and trajectory CSVs are present, but the GPU-extracted golden

```

```

features + `.rallies.csv` labels are not yet here):
53     1. **[owner / GPU box]** `python -m backend.eval.fusion_golden --manifest
output/human_lovo_manifest.json --out-dir "%USERPROFILE%\OneDrive\rally-annotator\golden-
data\features\2026-06-15-n6" ? produces `_golden_features.json` for the 6 golden videos. (GPU
+ WSL + videos + trajectory CSVs + `.rallies.csv` labels are owner-side only.)
54     2. Artifacts land in the **shared cloud folder** `C:\Users\avidu\OneDrive\rally-
annotator\golden-data\` (OneDrive auto-syncs both ways). Convention + workflow:
**[`GOLDEN_DATA_SHARING.md`](data_pipeline/GOLDEN_DATA_SHARING.md)**; tracked in **#175**.
55     3. **[CPU dev/agent session]** `python -m backend.eval.fusion_compare --features-dir
<shared>\features\2026-06-15-n6 --record` + `python -m backend.eval.ablation --features-dir
<shared>\... --granularity group` ? real person-feature ? F1; copy the JSONs into `output/` to
also pass the `skipif` litmus tests (`tests/test_ablation.py::test_litmus_reproduces_gate_a`).
56
57     **Also queued (owner-gated / pending input):**
58     - Ship-gate target calibration ? **#172** (owner decision; needs corpus growth).
59     - Tests reorganization ? blocked on the owner's grouping metric (provide it ? an agent
will do it).
60     - Restore CI ? the Actions outage is account-level (billing/minutes), not a repo/config
bug.
61
62     **Leaving (bulky / M2 ? no start without owner go):** ActionFormer (M0.4), event-index
DB migration (M0.1), Gemini-silver purge (M0.3), cost-to-Gen-2 benchmark + ByteTrack
deprecation, multisport, repo rename, PMF execution.
63
64     ## Where we are (one paragraph)
65     ADR-008 ratified ("repo split driven by productionization, not isolation"): training
lives in-repo (`training/`)
66     behind a CI-enforced import wall; the featurizer is single-sourced
(`backend/features/rally_seq.py`, train==serve);
67     the **Gen-0 harness ships** (`python -m training.gen0.harness --manifest
output/human_lovo_manifest.json --floor
68     eval_baselines/heuristic_lovo_n6.json --version <semver>`) and produced artifact v0.1.0
(LOVO **0.626**, floor passed,
69     sign test 5/6 wins p=0.109 ? honestly not significant, ship=False). The **n=6 owned
corpus is in the collector**
70     (262 rallies, Proprietary, commercial export = owned data only after the ShuttleSet
reclassification). The **Gemini
71     battery is done** (see table) ? the moat is quantified.
72
73     ## The quality table that now drives strategy (IoU 0.5 vs human golden labels)
74     | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |
75     |---|---|---|
76     | Heuristic (velocity windowing) | 0.657 | 0.157 |
77     | **Gen-0 owned head** (LOVO, n=6) | 0.744 | 0.561 |
78     | Two-stage WASB+Gemini refine | 0.83 | ? |
79     | **Gemini wrapper** (`gemini-flash-latest`) | **0.93** | **0.66** |
80
81     **Reading:** clean footage is commoditized (serve it with Gemini for cents). **Multi-
court drops the commodity
82     wrapper to 0.66 ? that is the owned model's ship target** (its FPs are background-game
pickups + dead-time merges,
83     the exact contrast-metric confound). Gen-0 is 0.10 behind with only n=6 ? a corpus-
growth-sized gap, not an
84     architecture-sized one. Two-stage refine = the cost-efficient serving path for LONG
tapes (uploads candidate clips
85     only); wrapper 503-flakiness (observed all session) makes the provider-fallback registry
load-bearing.
86     Baselines tracked in `eval_baselines/` (heuristic_lovo_n6.json,
gemini_wrapper_2026-06-10.json).
87
88     ## Rally-detection quality track (design merged 2026-06-13, owner-gated; implementation
in progress)
89     **This is the track a desktop/local agent is slated to pick up.** Design docs landed (PR
#112 + session). All implementation is **owner-gated** (per explicit in the plan); each step =
ONE PR off `master`.

```

90

91 **Hardening loop (T1-T5 positive tests + audit + finale) COMPLETE (2026-06-14, PRs #156/#157/#160/#161/#163 + #165 + final records).** See archived `\docs\archives\CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md` (Living Log with every pre/post/review) and `\docs\archives\HARDENING_LOOP_HANDOFF-COMPLETED-2026-06-14.md` (full rules + STATE). Top-level now pointer stubs per #165 review. All 5 items + verification + archive done; clean slate.

92

93 - **Docs:** `\docs\MULTI_SIGNAL_FUSION_PLAN.md` (shuttle + person + audio, one fusion layer) .

94 `\docs\EXPERIMENT_HARNESS.md` (A/B + shadow + promotion gate, zero-regression) .

95 `\docs\ROORA_LABELING_GUIDELINES.md` (v8) . `\docs\HOW_RALLY_DETECTION_WORKS.md` (2-layer mental model).

96 - **Key finding (no re-investigation needed):** the owner's "held-shuttle counted as a rally" bug is **NOT a missing capability** ? `\backend\pipeline\detectors\rally_gate.py` already implements the net-crossing `**"Gate A"*` (`\apply_gate(..., min_crossings=2)`) that kills service-feed/held-shuttle windows, but it is `**only called from eval drivers**` (and there was no `\min_crossings` knob in `\IndexingConfig`). `**Wiring Gate A into the live segmenter (now via fusion_hybrid)` was the single cheapest, highest-confidence quality win.** (P2 FusionSegmenter specifics: registry `\fusion_hybrid`, `**default-OFF**`, seeds from substrate, persists `\net_crossings` + optional person features, optional served Gate A/B via config knobs; adversarially reviewed; suite green. See `\tests\test_served_gate_a_regression.py` for the honest finding that on production windowing Gate A is neutral/negative in some cases, hence kept opt-in with knob for safety, and the leaderboard revert note.)

97 - `**Current status (2026-06-14):**` P1 (person cue extraction to reusable `\PersonDetector`) + P2 (FusionSegmenter + Gate A signal + served handoff gating) + experiment harness P1 DONE (adversarial review + NO-REGRESSION; suite green). `**P3 next (this PR + follow-ups):**` learned fusion (person features ? harness `\compare` LOVO lift on the 6 golden videos; eager-person-compute optimisation); then P4 audio via `hit_detector`.

98 - `**Pick-up order (each = ONE PR off \master, owner approval first, flag-gated default-OFF, promote only on measured LOVO):**`

99 1. (P2 first-step, largely done via `fusion_hybrid`) Wire Gate A into production served path + config knob (`min_crossings`); apply in hybrids at runtime. (Now available opt-in; served neutral on production windowing per litmus ? kept opt-in with knob for safety.)

100 2. P3 first step (this work item): harness `\compare` LOVO litmus for person-on variant (`players_in_court` + `two_sided` + `colocation` etc. from `\fusion_features`) on the 6 golden; add the eager compute optimisation note/landing. Pure + harness (testable here).

101 3. P3/P4 continuation: audio cue integration + full learned promotion if lift proven.

102 - `**Discipline (from EXPERIMENT_HARNESS.md):**` every new cue `**flag-gated, default OFF**`; promote to default `**only on measured LOVO lift**` through `\run_regression.py` (exit 0/1/3); pinned `\CONTROL` ? rollback = config flip. Harness re-uses ~80% existing code.

103 - ? Gate A, person cue logic, harness, and pure fusion paths are testable here; the heavier real-video / golden LOVO / ablation confirmation waits on the GPU-extracted golden features + `\.rallies.csv` labels, which are not yet present in this checkout.

104 - `**External due-diligence review inputs (2026-06-17) ? PICK UP AFTER` current ideas are exhausted; full details + citations in `\RALLY_QUALITY_RESEARCH.md` (`RALLY_QUALITY_RESEARCH.md`) §5.6.** An owner-commissioned 31-reference review of the quality report `*validated*` the methodology and surfaced these `**future**` levers (research inputs, not active work, not new claims; verify citations before any submission):

105 1. `**"First-mile" paper framing (highest-value):**` stroke-classifiers (BST/DSTA) + the big datasets (ShuttleSet ~36k strokes, Fine-Badminton ~27k actions) all `**assume a pre-trimmed rally clip already exists**` ? we solve the ignored prerequisite. Pitch the paper as `**the missing untrimmed-video preprocessing pipeline that lets BST/MonoTrack run on raw footage.**`

106 2. `**Ground R2 in Action-Spotting / Precise-Event-Spotting**` (SoccerNet/TennisSet/T-DEED) ? turns R2 into a defensible publishable contribution.

107 3. `**Add domain baselines for any submission:**` MonoTrack + BST on ShuttleSet / Fine-Badminton (we have only heuristic + Gemini today).

108 4. `**Prefer TemporalMaxer over ActionFormer**` for the M0.4 scorer swap (parameter-free max-pool, ~2.8x fewer GMACs, small-N robust) ? see "M0.4 next increment" below.

109 5. `**Court polygons ? pose ? the rally-END lever:**` masking on-court players can resolve `*occluded end-of-rally states*` (the #1 remaining gap vs Gemini), not just spectator noise.

110 6. `**Name the eval?serve lesson "pipeline distribution skew"***` ? a citable paper "lesson."

111

```

112     **Discipline reminder:** owner approval before starting any owner-gated item; one PR per
step; babysit the PR (poll/fix/reply) until merge observed, *then* advance.
113
114     ## Prioritized next steps
115     1. **[owner, in progress] GROW THE GOLDEN CORPUS** ? the single highest-leverage move;
every lever below feeds on it.
116         Priority order for new videos: **(a) multi-court badminton** (the target distribution
AND where the ship target
117         lives), **(b) ?3 matches per new sport ? table tennis first** (WASB transfers at F1
0.909; pickleball labels are
118         corpus-gold but not trainable until a clean MIT/Apache ball detector exists), (c)
capture variety per
119         `VIDEO_RECORDING_GUIDELINES.md`; **adult sessions only for the training pool**
(consent guardrail). Per video:
120         label ? `curate import-local --normalize --license Proprietary --fps <true fps>` ?
re-run the Gen-0 harness.
121         The paired sign test needs n: at n=10 with 9 wins, p?0.011 ? significance is
reachable this month.
122     2. **[owner decision] Set the ship-gate target** ? now informed: floor = heuristic 0.480
(necessary), **multi-court
123         reference = wrapper 0.66** (the number that makes the owned model worth serving),
clean-footage ceiling = 0.93
124         (not the target ? Gemini serves that tier). Suggested shape: "LOVO F1@tIoU0.5 ? 0.70
on multi-court folds with
125         paired-sign p<0.05 vs heuristic" ? owner to ratify/adjust.
126     3. **M0.4 next increment ($0/local):** swap the Gen-0 LogReg for a lightweight temporal
head (MIT, minutes on the
127         2070S ? NOT the heavyweight Gen-2 stack) inside the existing harness; per-video
threshold calibration (the
128         LOVO-vs-oracle gap is visible per fold in the harness output). **Prefer
TemporalMaxer** (parameter-free
129         max-pool, ~2.8x fewer GMACS, small-N robust) over ActionFormer/ASFormer per the
2026-06-17 external review
130         (`RALLY_QUALITY_RESEARCH.md` $5.6) ? better fit for the cheap/local/fast mandate.
131     4. **Multisport thread:** ingest Roora pickleball/TT CSVs (`import-local --normalize`);
merged-prototype LOVO across
132         badminton+TT once TT has ?3 matches; identify a clean pickleball ball detector. ? TT-
on-WASB serving stays gated
133         by C3 (provenance multiplies per sport).
134     5. **PMF validation (cheaper than any engineering month):** C13 academy interviews (+
the two added questions:
135         "multi-court footage nobody watches?" / "pay per-match to index it?"); camera pilot
at the warmest 2-3 academies
136         (consent at capture; demo served via the Gemini path or human-polish ? the reel
pattern exists:
137         `output/MahadevpuraSingles_highlights.mp4`). **Field kit COMPLETE (2026-06-12, owner
ratified the
138         physical-visit approach; reframed product-feedback-first per owner spec):**
139         `docs/PRODUCT_FIELD_ASSOCIATE_FLYER.md` (the recruiting flyer ? owner fills ?/contact
placeholders and recruits) +
140         `docs/FIELD_VISIT_PLAYBOOK.md` (the hired associate's visit playbook: coach + student
script, demo-after-Q6
141         rule, ground rules, Phase-2 record-and-show-back stretch goal) +
`docs/FIELD_VISIT_ANSWER_SHEET.md` (per-visit
142         capture form incl. demo-reaction + student blocks) + `docs/PMF_FIELD_SIGNALS.md`
(G/A/R anchors + **pre-registered
143         decision rules** ? PROCEED / PIVOT A-C / FALSIFIED thresholds fixed before the first
visit).
144     6. **Housekeeping:** ~~rotate the Gemini API key~~ **DONE 2026-06-10** (owner-verified
working; confirm old-key
145         deletion in AI Studio); **rename repo off "badminton"*** (owner action,
146         in BACKLOG); collector=SoR confirmation; off-peak niceties: none pending ? the
battery completed.
147
148     ## Pointers / artifacts

```

```

149 - Roadmap: `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md` · Topology: ADR-008
(`docs/archives/decisions/DECISIONS.md`) ·
150 Study: `docs/archives/research/OWNED_MODEL_TRAINING_STUDY.md` · Quality history:
151 `docs/archives/quality-iterations/2026-06-multivideo-lovo/`.
152 - Harness: `training/gen0/harness.py` (one command ? train ? LOVO ? gate ? artifact
bundle; weights gitignored,
153 manifest reviewable). Featurizer: `backend/features/rally_seq.py` (FEATURE_SCHEMA v1).
154 - Corpus: collector DB (system-of-record, 6 matches/262 rallies) · eval manifest
`output/human_lovo_manifest.json`
155 (gitignored) · durable trajectories + Roora files: `C:\Users\avidu\Projects\Annotation
Setup\Collect\Trajectories\` ·
156 labels: `...\Collect\Golden Labelled\`.
157 - Gemini ops learnings: serial uploads (`indexing.ai_max_workers=1`) + chunk re-encode
158 (`indexing.ai_chunk_scale_width=1280`) + 6-attempt exponential generate retry ? a
cold/prepaid project burst-throttles
159 with a MISLEADING "credits depleted" 429; ramp gradually.
160 - ? No owned-model implementation beyond greenlit M0.x until owner greenlights the next
milestone. Committed next
161 PLATFORM slice = async job model (single-tenant).
162

```

7. user (2026-06-21T13:56:24.455Z)

```

1 # Rally-Detection Quality: Research, Eval Framework & Project Plan
2
3 **Status:** Foundation doc (2026-06-16), run as a **tracked project** (§7) ? "Bridge
the
4 local?Gemini rally-detection gap,"* **completed when all committed tiers land and a
single quality
5 number is attributed (§7.1).** The empirical + research backbone for closing the gap
between the
6 **local, no-cloud rally segmenter** and the **Gemini** path. Iterations land as their
own measured
7 PRs, logged in [`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md); this doc is the plan
they execute against.
8
9 **How to read this:** §1 frames the gap with measured numbers. §2 is the precise, code-
grounded
10 diagnosis (read before proposing fixes). §3 inventories the eval/experiment harness we
already
11 have. §4 specifies how we strengthen that harness (the "really strong evaluation
framework").
12 §5 is the cited external-research synthesis (4 dimensions). §6 is the prioritized
roadmap with
13 falsifiable success criteria. **§7 is the project plan ? work-item tracker, sequencing,
the single
14 number we attribute to the effort, and the definition of done.** §8 recaps the non-
negotiable
15 constraints + cross-links.
16
17 > **This doc does NOT duplicate prior design.** It builds on, and cross-links:
18 > [`HOW_RALLY_DETECTION_WORKS.md`](HOW_RALLY_DETECTION_WORKS.md) (the 2-layer mental
model),
19 > [`EXPERIMENT_HARNESS.md`](EXPERIMENT_HARNESS.md) (the A/B + LOVO harness, P1 shipped),
20 > [`MULTI_SIGNAL_FUSION_PLAN.md`](MULTI_SIGNAL_FUSION_PLAN.md) (shuttle+person+audio
fusion),
21 > [`ANALYZERS_AND_RECONCILERS.md`](ANALYZERS_AND_RECONCILERS.md) (the signal-fusion
framework +
22 > §13 literature review ? **already** establishes that our skeleton is standard TAL/TAS
+
23 > late-fusion; we *cite, don't re-derive*), [`ABLATION_LAYER.md`](ABLATION_LAYER.md),
24 > [`EVALUATION.md`](EVALUATION.md),
[ `OWNED_MODEL_IMPLEMENTATION_PLAN.md` ](OWNED_MODEL_IMPLEMENTATION_PLAN.md)
25 > (TAL-head-first ? VLM), and [`GOLDEN_SET_VENDORING.md`](archives/roora-
vendor/GOLDEN_SET_VENDORING.md) (corpus growth).

```

```

26
27 ---
28
29 ## 1. The gap (measured)
30
31 Two paths produce rally segments today. Gemini (cloud VLM, watches the video
semantically)
32 is strong; the local path (WASB shuttle trajectory ? velocity windowing, no cloud)
lags
33 badly. Measured F1 vs golden labels ([`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md)):
34
35 | Path | Mahadevpura (clean, single court) | Kushagra (multi-court) |
36 |---|---|---|
37 | Heuristic velocity windowing (local today) | 0.657 | 0.157 |
38 | Gen-0 owned head (LogReg on 5 trajectory features, LOVO n=6) | 0.744 | 0.561 |
39 | Two-stage WASB + Gemini refine | 0.83 | ? |
40 | Gemini wrapper (gemini-flash) | 0.93 | 0.66 |
41
42 The local heuristic is footage-fragile: fine on a clean single court, collapses on
43 multi-court / busy venues. The owned Gen-0 head already closes most of the multi-court
gap
44 (0.561 vs 0.66) ? the residual is corpus-sized, not architecture-sized.
45
46 ### 1.1 The 2026-06-16 local-vs-Gemini divergence run (no golden labels)
47
48 A fully-local run (`wasb_hybrid --skip-ai-handoff`) vs Gemini on 3 fresh matches
49 (Kushagra/Yash June-12; report: `?/June 12 Local-vs-Gemini/_LOCAL_vs_GEMINI_REPORT.md`,
50 generator `scratch/compare_local_vs_gemini.py`). Gemini is a strong reference here,
NOT
51 ground truth (these clips are unlabeled) ? so this is an agreement/divergence study:
52
53 | | Gemini | Local (WASB no-AI) |
54 |---|---|---|
55 | rallies / windows (3 videos) | 59 | 12 |
56 | total "play" | 6:28 | 17:41 |
57 | mean window duration | 6.1 s | 65 s (10.6x) |
58 | true misses (gemini-only) | ? | 0 |
59 | merge groups (one local window ? 2 Gemini rallies) | ? | 7 |
60
61 Headline: the local path does not miss rallies ? it fails to segment them. One
62 mahadevpura-1 window spanned the entire 174 s clip. Local "play" is 2.7x Gemini's
because its
63 mega-windows swallow the dead time between rallies.
64
65 ---
66
67 ## 2. Diagnosis (code-grounded ? read before proposing fixes)
68
69 The local pipeline is two layers (see
[`HOW_RALLY_DETECTION_WORKS.md`](HOW_RALLY_DETECTION_WORKS.md)):
70 (L1) detection ? per-frame shuttle position (WASB/TrackNet); and (L2)
segmentation ?
71 turn the trajectory into rally intervals. The gap is almost entirely L2.
72
73 ### 2.1 L1 (shuttle tracking) is NOT the bottleneck
74 On the divergence clips WASB reported the shuttle visible in 96?99 % of frames
(mahadevpura-1
75 98.6 %, GX020125 72.8 %, mahadevpura-2 78.8 %). The trajectory signal is present and
dense.
76
77 ### 2.2 L2 (segmentation) is the bottleneck ? and the flaw is conceptual
78 The windowing is `trajectory_to_action_windows`
79 ([`backend/pipeline/detectors/tracknet_runner.py:256`](../backend/pipeline/detectors/tra
cknet_runner.py)):
80

```



```

81 - A frame is "active" iff the shuttle is visible AND its per-second-normalized
speed**
82 exceeds `velocity_threshold` (default 0.02). Velocity = `(dist / frame_width) * (fps /
dframes)`
83 ? already fps-normalized (? fraction of frame-width per second).
84 - Active frames within `dead_time_merge_gap` (2.0 s) are merged into one window.
85 - Windows shorter than `min_action_window_duration` (2.0 s) are dropped.
86
87 The conceptual flaw: "the shuttle is moving" is treated as "a rally is in progress."**
During
88 dead time the shuttle is still visible and gently moving (being picked up, walked back,
knock-ups,
89 slow serve prep) ? those frames read as active. The 2 s merge then bridges the brief
gaps, and
90 because (at the time this was diagnosed) there was no `max_window_duration` cap and no
91 inactivity/boundary split**, a single window grew until the shuttle finally went
invisible ? which,
92 at 96?99 % visibility, is rarely. Result: rally + dead-time + rally + ? collapse into a
few
93 mega-windows. *(This is the failure W1/R4 fix; both are now default-ON since the R1
milestone, PR
94 #204 ? `max_window_duration=6`, `window_inactivity_end=1.5` in `IndexingConfig`.)*
95
96 > Correction to an earlier hypothesis (honesty note). A first pass guessed the cause
was
97 > "velocity threshold not fps-invariant." It IS fps-invariant (the `fps/dframes`
factor). The
98 > real cause is the *motion-equals-rally* conflation + 2 s bridging + no upper bound.
fps matters
99 > only secondarily: denser 60 fps sampling keeps more slow-motion dead-time frames above
the
100 > speed threshold. The `_LOCAL_vs_GEMINI_REPORT.md` has been corrected to match.
101
102 ### 2.3 What this implies
103 The local path has no rally-STATE signal ? only "is the shuttle moving." Bridging
the gap is
104 fundamentally about giving L2 signals that separate *rally* from *shuttle-merely-moving*
105 (net-crossings, two-sided play, sustained fast exchange, hit sounds) and a segmenter
that
106 produces *bounded, well-cut* intervals ? exactly the multi-signal-fusion + learned-
segmenter
107 direction the existing docs lay out. This doc makes that concrete and measurable.
108
109 ---
110
111 ## 3. What we already have (eval/experiment harness inventory)
112
113 The harness is rich and composable ? we extend it, not reinvent it. (Modules under
114 `backend/eval/`.)
115
116 **Scoring & matching**
117 - `metrics.py` ? temporal-IoU interval matching; precision/recall/F1, mean_iou, boundary
errors.
118 - `segmentation_metrics.py` ? F1@{0.1,0.25,0.5}, mAP@tIoU (0.3?0.7),
segment_count_ratio
119 (the over/under-segmentation diagnostic). (Edit/segmental score deliberately omitted
to date.)
120
121 **Statistical rigor (small-n)**
122 - `eval_stats.py` ? bootstrap 95 % CIs, paired exact sign-test, `paired_delta_ci`,
123 `grouped_folds` (LOGO), `out_of_fold_predictions` (anti-stacking-leakage). The
promotion bar =
124 "F1 CI excludes 0 AND sign-test agrees."
125 - `score_calibration.py` ? Platt + isotonic calibrators, Brier + ECE.
126

```

```

127     **Variants, A/B, leaderboard**
128     - `experiment.py` ? `Variant` (declarative recipe: segmenter + config + cue_flags +
gates +
129         learned-classifier + threshold/merge_gap/calibrate), `compare()` (LOVO-honest A/B with
an
130         "honest" stats block), `compare_unlabeled()` (agreement on unlabeled footage),
`CONTROL`
131         (pinned baseline), `record_experiment()` ?
`eval_baselines/experiment_leaderboard.json`.
132     - `ablation.py` ? leave-one-signal-out marginal ?F1 with CI + sign-test + a no-over-
claim verdict
133         vocabulary (`earns_keep` / `suggestive` / `inert` / `structural_protected`).
134     - `classifier.py` ? dependency-free numpy `LogisticRegression` (the learned scorer;
JSON-serializable).
135
136     **Eval?serve discipline + gates**
137     - `windowing.py` ? single-sources the **SERVED** preset from `IndexingConfig` defaults
so eval
138         and production can't silently diverge; `PROPOSAL` (recall-oriented) is where the
golden feature
139         substrate is built. (This exists *because* Gate-A measured **+0.074 on PROPOSAL but
?0.008 on
140         SERVED** and was reverted ? the canonical eval?serve lesson.)
141     - `promotion.py` / `served_gate_a.py` ? promote a cue only if it does NOT regress at the
served
142         windowing; `per_user_gate.py` / `per_user_regression.py` ? per-user analogues;
`regression.py` ?
143         golden-set no-regression baseline with a GT-hash to detect label changes.
144
145     **Golden substrate**
146     - `gt_loader.py` ? one canonical GT loader (both legacy formats). `fusion_golden.py` /
147         `fusion_compare.py` / `fusion_audio.py` ? extract replayable per-window feature JSONs
148         (shuttle 5 + person 5 + audio 3) so ablation/compare re-score on CPU without re-
detecting.
149     - **Corpus today: 15 golden videos** (started at the 6-video label-set hash
`50d98ffbc370`, since
150         grown with the GX/Boxhill multicourt clips); more expected via the
151         Roora vendoring track ([`GOLDEN_SET_VENDORING.md`](archives/roora-
vendor/GOLDEN_SET_VENDORING.md)).
152
153     **Signals available to L2 (built, default-OFF):** net-crossing count (`rally_gate.py`,
Gate-A),
154     5 person/court features (`person_detector.py` + `fusion_features.py`), 3 audio-hit
features
155     (`audio_features.py`). Fusion hooks: `fusion_hybrid.py::_enrich_candidate_stats` /
156     `_select_for_handoff`. Persisted in `candidate_segments.stats`.
157
158     ### 3.1 Known gaps in the harness (what $4 fills)
159     1. **Over-segmentation is under-surfaced in the default report.** `segment_count_ratio`
exists but
160         the merge/split *breakdown* (the very thing that exposed the over-merge) lives only
in a scratch
161         script. tIoU-F1 alone is **blind to over-merge** ? a giant window can still "match"
via overlap.
162     2. **No segmental edit-score** (Levenshtein over the predicted/gt label sequence) ? the
standard
163         TAS over-segmentation metric.
164     3. **Per-cue calibration is fitted but never *measured* post-hoc** (no ECE/Brier delta
reported).
165     4. **`out_of_fold_predictions` (anti-leakage stacking) is built but not wired to an
arbiter.**
166     5. **No active-learning loop** ? nothing tells us *which* unlabeled video to label next.
167     6. **No single "rally-segmentation eval" entrypoint** ? the pieces exist but a one-
command
168         `score this segmenter vs golden, honestly` runner would make every iteration cheap.

```

```

169      7. **The local-vs-Gemini agreement track is ad-hoc** (a scratch script), not a first-
class
170      *shadow* signal alongside the golden gate.
171
172      ---
173
174      ## 4. The strengthened evaluation framework (the "really strong harness")
175
176      Design principles (inherited, non-negotiable):
177      - **Signals ? learned scorer, NO special-casing**
([`ANALYZERS_AND_RECONCILERS.md`](ANALYZERS_AND_RECONCILERS.md) §1):
178      every new cue is a feature column the scorer learns to weight; never an `if/else` in
the decision path.
179      - **Promote only on measured lift; default OFF; revert is one flag**
([`EXPERIMENT_HARNESS.md`](EXPERIMENT_HARNESS.md) §6).
180      - **eval ? serve:** a harness win must be re-confirmed on the *served* windowing before
any default flip.
181      - **Honest small-n stats:** never a bare mean at  $n \geq 16$ ; always CI + paired sign-test
(C1).
182
183      ### 4.1 Metric set the leaderboard should standardize on
184      Surface ALL of these in the default `compare()` report (most already computed; the task
is to
185      *report* them together so over-merge can never hide again):
186
187      - **tIoU-F1 @0.5** (primary) **and F1@{0.1,0.25,0.5}** (segmental, TAS-standard).
188      - **mAP@tIoU** (0.3?0.7) when predictions carry confidences.
189      - **`segment_count_ratio`** + a **merge-rate / split-rate** breakdown (fraction of GT
rallies that
190      share a predicted window = under-segmentation; fraction of predicted windows that span
?2 GT =
191      the over-merge signal). Promote the scratch `compare_local_vs_gemini.py` matching into
a reusable
192      `backend/eval/` helper.
193      - **Segmental edit (Levenshtein) score** ? the canonical over-segmentation metric we
currently lack.
194      - **Boundary timing** (mean start/end error on matched pairs) ? already in `metrics.py`,
surface it.
195      - **Calibration (ECE/Brier)** reported pre/post per-cue calibration (close gap #3).
196
197      > **Why:** plain tIoU-F1 rewarded our mega-windows (they "overlap" real rallies). The
metric set
198      > above makes over/under-segmentation *first-class*, so a fix that improves segmentation
is
199      > visible even when tIoU-F1 barely moves ? and vice-versa.
200
201      ### 4.2 Splits & significance
202      - **Leave-one-video-out, grouped by match** (never leave-one-*window*-out ? windows
within a video
203      are correlated). Keep `grouped_folds`. Note the **RL00CV bias** caveat for tiny n (see
§5.4).
204      - **Bootstrap CI + paired exact sign-test** on per-video held-out F1 (C1). A result is
*promotable*
205      only if CI excludes 0 **and** the sign-test agrees, **and** it survives the served-
windowing check.
206      - **Anti-leakage:** per-LOVO-fold threshold tuning + calibration on the TRAIN fold only
(already
207      enforced in `experiment._calibrate_and_tune`); wire `out_of_fold_predictions` once an
arbiter exists.
208
209      ### 4.3 The golden corpus + active learning (the highest-leverage lever)
210      The single biggest quality lever is **more labeled matches**
([`NEXT_STEPS.md`](NEXT_STEPS.md)):
211      "corpus growth = highest leverage". With ~6 ? ~16 videos:
212      - Grow via the Roorra vendoring flow ([`GOLDEN_SET_VENDURING.md`](archives/roora-

```

```

vendor/GOLDEN_SET_VENDORING.md),
213     [``ROORA_LABELING_GUIDELINES.md``](archives/roora-vendor/ROORA_LABELING_GUIDELINES.md))
? court calibration, IAA, the
214     3 non-negotiables (ignore dead time, every rally, one-contact = one shot).
215     - **Active learning:** stop labeling videos at random. Score the *unlabeled* corpus and
label the
216     videos where the local segmenter is most *uncertain* / most *divergent from Gemini* /
most
217     *diverse* (new venue, fps, lighting). §5.3 surveys the methods; the local-vs-Gemini
agreement
218     signal (§4.4) is a ready proxy for "where do we disagree most ? label that."
219
220     ### 4.4 Shadow evaluation: Gemini-agreement as a free, unlimited signal (NOT a gate, NOT
training data)
221     We have exactly one strong reference available on *unlabeled* footage at inference time:
**Gemini**.
222     - `experiment.compare_unlabeled()` already scores agreement (agreed / a-only / b-only)
without labels.
223     - Use it as a **shadow metric**: run any candidate segmenter on a large unlabeled pool,
measure
224     agreement-with-Gemini, and watch it as a *trend* ? cheap, unlimited, catches
regressions the 6
225     golden videos can't. **It is a shadow signal only.** Golden labels remain the gate;
Gemini is not
226     ground truth (it has its own errors) and ? per the licensing guardrail (§8) ? **its
outputs may
227     never become training data for owned weights.**
228
229     ### 4.5 One-command rally-segmentation eval (close gap #6)
230     Wrap the above into a single reusable entrypoint (a future PR ? not built here): given a
segmenter
231     variant + the golden manifest, emit the full §4.1 metric set + the honest stats block +
a leaderboard
232     row, and (optionally) the §4.4 shadow agreement on an unlabeled pool. Every iteration in
§6 is then
233     "run the entrypoint, read the honest delta, promote-or-not."
234
235     ---
236
237     ## 5. External research synthesis (cited, adversarially verified)
238
239     > Produced by a deep-research run (6 search angles ? 27 sources ? 126 candidate claims ?
25
240     > 3-vote-adversarially-verified, 23 confirmed / 2 refuted). **Cite, don't over-claim.**
Two claims
241     > were *killed* in verification ? recorded in §5.5 so we don't repeat them. **License
caveat:** the
242     > code repos below are research releases; **independently confirm each repo's actual
license is
243     > MIT/Apache-2.0/BSD before vendoring any of it** (this survey did not verify every
license file).
244     > **Governance:** none of these techniques require training on paid-LLM outputs ? all
are compatible
245     > with our no-Gemini-training-data guardrail (§8).
246
247     ### 5.1 Windowing & boundaries (TAL/TAS) ? the explicit-boundary machinery we lack
248     The recurring lesson: modern action-localization models have an **explicit boundary
mechanism**;
249     our heuristic has none (it merges on a velocity+gap rule). The transferable *idea* is
"predict a
250     boundary and cut there," but the deep models are built for **hundreds+ of training
videos**, so at
251     n=6?16 they are **research bets**, not drop-ins.
252
253     - **ASRF ? Action Segment Refinement Framework** (Ishikawa et al., WACV 2021,

```

[arXiv:2007.06866](https://arxiv.org/abs/2007.06866)).

254 A shared feature extractor feeds **two branches**: frame-wise classification + a
Boundary

255 Regression Branch; predicted boundaries refine/split the segmentation (+13.7% edit,
+16.1%

256 segmental F1 over MS-TCN-era SOTA on 50Salads/GTEA/Breakfast). **Why it fits:** the
boundary head

257 is exactly the missing "cut here" mechanism. **? Two caveats:** (1) ASRF's headline
goal is to

258 **reduce** over-segmentation (merge spurious fragments) ? the opposite of our
over-merge; only the

259 **split-at-a-detected-boundary** sub-mechanism transfers. (2) The claim that ASRF's
boundary branch

260 is a **model-independent post-hoc splitter** you can bolt onto any dense signal was
REFUTED

261 (\$5.5) ? you must train the two-branch model, not graft the branch. **Effort/payoff:**
research bet.

262

263 - **TriDet** (Shi et al., CVPR 2023,
[arXiv:2303.07347](https://arxiv.org/abs/2303.07347)). A

264 **Trident-head** models each boundary as a **relative probability distribution** over
local bins

265 (boundary = expected value over neighboring instants) ? built for imprecise/ambiguous
boundaries.

266 69.3% avg mAP on THUMOS14 (> ActionFormer 66.8% at ~75% latency), **convolution-only**
(SGP), no

267 self-attention (attractive on a single local GPU). Ablation shows the Trident-head's
gain

268 concentrates at strict tIoU=0.7 ? it buys **boundary precision**. **Effort/payoff:**
research bet

269 (trained on 100s of videos); the portable piece is the relative-boundary idea.

270

271 - **ActionFormer** (ECCV 2022, [arXiv:2202.07925](https://arxiv.org/abs/2202.07925)).
The reference

272 single-stage anchor-free template: multiscale pyramid + local self-attention; classify
every

273 moment + regress its boundaries (71.0% mAP@0.5 THUMOS14). This is the family
TriDet/TemporalMaxer/

274 CausalTAD extend, and the natural target if we ever graduate from heuristic+LogReg.
Research bet.

275

276 - **TemporalMaxer** (CVPRW 2023, [arXiv:2303.09055](https://arxiv.org/abs/2303.09055)) ?
the most

277 important small-data signal in this dimension. It replaces ActionFormer's self-
attention block

278 with a **parameter-free max-pool** and **beats** it (67.7 vs 66.8 avg mAP) at **~4x**
fewer params.

279 The ablation is the takeaway: a **learned 1-D conv block** is the **WORST** (59.4,
overfits redundant

280 features); parameter-free pooling wins. **Why it fits:** direct evidence to **keep our**
scorer

281 **parameter-light** (the numpy LogisticRegression over a handful of features is the
right choice at

282 our scale, not a placeholder to rush past). **Effort/payoff:** a discipline (free) ?
adopt the

283 principle, not the model. (Caveat: the overfitting evidence is about feature
redundancy on a

284 mid-size set, a directional inference to our small-n setting.)*

285

286 - **CausalTAD** (2024, [arXiv:2407.17792](https://arxiv.org/abs/2407.17792)) ?
steering-away note.

287 Causal-attention masking (attend only to past/present) won EPIC-Kitchens/Ego4D 2024
tracks, but it

288 adds **no** boundary-split module. Useful later for **online/streaming** rally
detection; **NOT** the

289 over-merge fix. Listed so we don't mistake causal attention for the remedy.

290

291 - **Metrics (the actionable cross-cut):** the field leads with **segmental Edit**
(Levenshtein) score

292 and **F1@k**, not frame accuracy (MoF) or plain tIoU ? see §5.4.

293

294 **### 5.2 Shuttle-trajectory ? rally signal processing ? published badminton systems**
already do what we lack

295 This is the highest-confidence, lowest-effort dimension: in-domain papers operationalize
the exact

296 boundary heuristics we're missing, in spirit license-compatible (re-confirm before
vending code).

297

298 - **See et al. ? Service & End-of-Rally Detection in Badminton** (2024 Springer chapter
299
[10.1007/978-3-031-69073-0_15](https://link.springer.com/chapter/10.1007/978-3-031-69073-0_15);
300 2025 SN Computer Science follow-on
[10.1007/s42979-025-03935-0](https://link.springer.com/article/10.1007/s42979-025-03935-0)).

301 ***(a) Rally END = a sustained-inactivity rule:** the rally ends when **~80%** of
consecutive frames

302 show no shuttle motion (**~95% accuracy** over 187 rallies). ***(b) Rally START**
(service) = a

303 learned classifier over court geometry + player poses (~26 features/frame × 12
frames ?

304 CNN/BiLSTM, **88.6%** accuracy). **Why it fits:** the inactivity rule is ***precisely***
the

305 end-of-rally cut our windowing lacks ? a few lines, directly breaks 65 s mega-windows.
?? Caveats:

306 the rule is a ***termination***, not a ***within-rally split***; results rest on paywalled
abstracts +

307 small test sets (35?187 rallies); the serve classifier needs the (default-OFF) person
cue + court

308 polygons. **Effort/payoff:** the inactivity rule = ***low-effort high-confidence win***;
serve

309 classifier = medium.

310

311 - **Hsu, Yu & Cheng ? trajectory + swing fusion** (Sensors 2024, 24(13):4372,
312 [10.3390/s24134372](https://www.mdpi.com/1424-8220/24/13/4372)). Fusing TrackNet
trajectory with a

313 YOLOv7 swing detector via a Shot Refinement Algorithm raises hit detection ***F1 72.3%**
? 90.5%*

314 (over 69 BWF matches / 1582 rallies). Hits are found by ***trajectory direction-**
reversal / peak

315 detection (**convert (x,y)?(frame,y) wave peaks**) then reconciled with the swing cue.
Why it fits:

316 validates our ***late/score-level fusion*** approach in-domain, and direction-
reversal/peak is a cheap

317 ***per-window feature*** marking contact events (a "rally has hit cadence" signal ? the
discriminator

318 a "shuttle-moving" heuristic lacks). **?? Important caveat (the one 2-1 vote):** this
paper does

319 ***within-rally stroke/hit* segmentation**, ***not*** rally start/end ? "it splits merged
rallies" is ***our**

320 extrapolation; frame it as "hit cadence plausibly helps locate rally boundaries," not
a proven

321 rally-splitter. **Effort/payoff:** the peak/direction-reversal feature = low-effort
win (feed the

322 scorer); full swing fusion = medium.

323

324 - **What best separates "rally" from "shuttle merely moving"** (synthesis): sustained-
inactivity

325 termination (See), hit cadence / direction-reversals (Hsu), net-crossings ?2 (our
Gate-A,

326 [`rally_gate.py`](../backend/pipeline/detectors/rally_gate.py)), two-sided court
occupancy (our

327 person features). None alone is sufficient; the spine (§4) says ****feed them all to the**
 learned
 328 scorer**.

329

330 **### 5.3 Small-data strategies ? most quality per label**
 331 - ****Timestamp / sparse supervision**** (Li et al., CVPR 2021,
 [arXiv:2103.06669](https://arxiv.org/abs/2103.06669);
 332 TAS survey Ding et al., [arXiv:2210.10352](https://arxiv.org/pdf/2210.10352)). Label
****one frame per**
 333 **action**** ? iteratively refine pseudo frame-labels ? ***comparable*** to full supervision
 (50Salads:
 334 timestamp 75.6% acc / 60.1 F1@50 vs full 83.7% / 70.1). ****Why it fits:**** sparse per-
 rally timestamps
 335 cost far less than dense boundary labels ? relevant if we ever train a frame-wise
 model.
 336 ****? Caveats:**** "comparable" hides a real ****~8?10 pt residual gap****; annotation savings
 are modest
 337 (~35%, annotators still watch the whole clip); benchmarks are cooking/egocentric ?
 transfer to 1-D
 338 shuttle trajectories is ****unproven****. ****Effort/payoff:**** research bet to pilot.
 339 - ****Parameter-light scorer**** ? see TemporalMaxer (§5.1): keep the numpy
 LogisticRegression; resist a
 340 learned temporal net until the corpus is much larger. ****Low-effort win (a**
 discipline).******
 341 - ****? OPEN (the research could NOT verify these ? honest gap):**** ***active learning***
 (which video/segment
 342 to label next via uncertainty/diversity/core-set) and the ***teacher-student /**
 distillation governance
 343 **boundary*** for using Gemini as a weak/pseudo-labeler returned ****no surviving verified**
 claim**. Our
 344 pragmatic, guardrail-safe stance (not from the literature ? from our own design): pick
 the next
 345 video to label by ****maximal local-vs-Gemini disagreement**** (the §4.4 shadow signal) +
****scorer**
 346 **uncertainty**** + ****venue/fps diversity****; and treat ****Gemini strictly as an inference-**
 time
 347 **weak-reference / eval oracle ? never a training-data source for owned weights**** (§8).
 These two
 348 topics deserve their own grounded research pass before we lean on them.
 349

350 **### 5.4 Evaluation rigor at small n ? lead with segmental metrics**
 351 - ****Frame accuracy (MoF) and plain tIoU are BLIND to over-segmentation**** (TAS survey,
 Ding et al.,
 352 [arXiv:2210.10352](https://arxiv.org/pdf/2210.10352); segmental-F1 origin Lea et al.,
 CVPR 2017).
 353 "The score can be high even when segments are fragmented." The over-merge that bit us
 is exactly
 354 what tIoU-F1 hides (a giant window still ***overlaps*** real rallies). ****Action:**** lead
 the leaderboard
 355 with the ****segmental Edit (Levenshtein) score**** (penalizes spurious
 ordering/fragmentation) ****and**
 356 **F1@{0.1,0.25,0.5}****, paired with our existing ``segment_count_ratio`` + the merge/split
 breakdown.
 357 ****We already have**** F1@k, mAP@tIoU, segment-count-ratio, bootstrap CIs, paired sign-
 test, Platt/
 358 isotonic + ECE/Brier ? the only ***missing*** primary metric is the ****Edit score**** (low-
 effort add).
 359 - ****? OPEN (not covered by a verified claim):**** IAA / label-noise level in the golden
 set, and which
 360 over-segmentation-robust loss (ASRF boundary/smoothing, MS-TCN truncated-MSE) best
 tolerates it at
 361 small n. Measure IAA as the corpus grows (Roora has an A/A overlap plan); revisit
 robust losses
 362 only when a frame-wise model is on the table.
 363

364 ### 5.5 Refuted claims (recorded so we do NOT repeat them)

365 1. **"ASRF's boundary branch is model-independent ? a post-hoc splitter applicable to any dense**

366 segmentation output."**** **REFUTED (0-3).**** Implication: we cannot bolt ASRF's boundary head onto

367 our trajectory signal as a drop-in; the two-branch model must be trained together. ?

368 *concept*, don't expect a plug-in.

369 2. **"A GFL-inspired Boundary Distribution Regression head is the mechanism most relevant to splitting**

370 merged windows."**** **REFUTED (0-3)**** (and the cited source had a suspect identifier). Do not adopt

371 this framing as established.

372

373 ### 5.6 External due-diligence review (2026-06-17) ? owner-commissioned; verify before submission

374 An owner-commissioned external review (31 cited sources) of the

375 RALLY_DETECTION_QUALITY_REPORT.md checkpoint ****validated the**

376 methodology****** (the decoupled 2-layer design, the R2 metric, LOVO + paired-sign-test discipline, the

377 consented-data + over-seg-aware-harness moat) and surfaced these ****borrowable inputs****. These are

378 *research inputs for later pickup* ? they add ****no measured numbers****, and the citations are the

379 review's own and ****must be re-verified before any submission****.

380

381 - ****Positioning ? the "first-mile" gap (highest-value).**** Badminton stroke-classification SOTA

382 (****BST**** ? Badminton Stroke-type Transformer, CVPR-W 2026; ****DSTA**** ? Decoupling Spatio-Temporal

383 Adapter) and the big datasets (****ShuttleSet**** ~36,492 strokes; ****Fine-Badminton**** ~27,597 actions)

384 all ****silently assume a pre-trimmed rally clip already exists**** ? they depend on a human having

385 segmented the untrimmed match. Our system solves that ignored ***prerequisite***. Reframe the paper as

386 **"the missing untrimmed-video preprocessing pipeline that lets BST/MonoTrack run on raw footage"** ?

387 a real, unoccupied gap and a far stronger narrative than "applied/systems."

388 - ****R2 ? Action-Spotting / Precise-Event-Spotting (AS/PES).**** Our asymmetric tolerance-match metric

389 sits at the ****TAL × AS intersection**** (interval task, tolerance-based match). Grounding it in the

390 AS/PES literature (SoccerNet, TenniSet, T-DEED, the PES survey [arXiv:2505.03991]) frames R2 as a

391 deliberate methodological contribution, not an ad-hoc metric ? strengthens §5.4 and the report's §9.

392 - ****Comparative baselines reviewers will require.**** Benchmark against ****MonoTrack**** and ****BST**** on

393 ****ShuttleSet / Fine-Badminton**** ? we currently have only a generic heuristic + the Gemini wrapper;

394 domain-specific baselines are a hard publication requirement.

395 - ****Scorer upgrade ? prefer TemporalMaxer.**** For the post-logreg scorer (corpus-gated; NEXT_STEPS M0.4),

396 ****TemporalMaxer**** (parameter-free temporal max-pool; ~2.8× fewer GMACs; robust at small-N,

397 [arXiv:2303.09055]) fits the cheap/local/fast mandate better than attention-heavy ActionFormer.

398 - ****Court polygons ? pose ? the rally-END lever.**** Court masking is not only the person-cue spectator

399 fix (§4 / MULTI_SIGNAL_FUSION_PLAN.md): once pose is isolated to on-court players it can resolve

400 ****ambiguous, occluded end-of-rally states**** ? feeding directly into the ****#1 remaining gap** (the


```

401 rally-END boundary)** , not just precision. Connect court-polygons to the R4/END-trim
work.
402 - **Name the eval?serve lesson "pipeline distribution skew."** Downstream-cue value
depends on the
403 upstream proposal distribution (the Gate-A episode) ? a citable framing and candidate
paper "lesson."
404 - **Citable Layer-1 anchor (public benchmark, NOT our footage).** WASB SBDT F1 95.6
(step=1) / 94.0
405 (step=3) vs MonoTrack 92.1, TrackNetV2 89.4, BallSeg 71.7 (WASB, BMVC 2023). Justifies
the frozen
406 detector choice; **keep distinct from our own, still-unmeasured footage domain-gap**
($2.1, a real
407 open limitation the review also flags).
408
409 **Sequencing:** these are *publication-prep + post-current-idea* levers. The
prerequisites remain
410 corpus growth to n?16 and **multi-court isolated as its own reported slice** ($6 Tier-3
/ $7) ? do
411 those first; the items above land once current windowing/fusion ideas are exhausted.
412
413 ---
414
415 ## 6. Prioritized roadmap (cheap ? durable)
416
417 Every item is **its own measured PR** : default-OFF, scored on the $4 metric set over the
golden set
418 with honest stats (CI + sign-test), **re-confirmed on the served windowing** (the
eval?serve gate),
419 logged in [ `QUALITY_ITERATIONS.md` ](QUALITY_ITERATIONS.md). Ordered by the research's
420 effort/confidence tiering. Each carries a **falsifiable success criterion** ? if it
isn't met, the
421 item does not promote (and that null is itself recorded).
422
423 ### Definition of "bridging the gap"
424 Local rally segmentation, measured **honestly LOVO-by-match on the golden set** ,
approaches the
425 Gemini reference (today **0.93 single-court / 0.66 multi-court** tIoU-F1) ? but tracked
primarily
426 through the **over-segmentation-sensitive metrics** (segmental Edit + F1@k + low
merge/split rate),
427 since those, not bare tIoU-F1, capture our actual failure. A secondary, free signal is
rising
428 **agreement-with-Gemini** on the unlabeled pool ($4.4, shadow only).
429
430 ### Tier 0 ? Make the failure measurable FIRST (do before any fix)
431 - **E1 · Over-segmentation-aware leaderboard.** Add the **segmental Edit (Levenshtein)
score** to
432 `segmentation_metrics.py` ; surface Edit + F1@{0.1,0.25,0.5} + `segment_count_ratio` +
a
433 **merge-rate / split-rate** breakdown (promote the
`scratch/compare_local_vs_gemini.py` matching
434 into a reusable `backend/eval/` helper) in the default `experiment.compare()` report;
add the
435 one-command rally-segmentation eval entrypoint ($4.5). **Success:** the current
heuristic's
436 over-merge shows up as a *number* (low Edit, high merge-rate) on the 6 golden videos,
and a
437 baseline row lands in the leaderboard. *(Low effort, high confidence ? research-backed
$5.4.)*
438
439 ### Tier 1 ? Low-effort, high-confidence wins
440 - **W1 · Bounded windowing (the direct over-merge fix).** Add a `max_window_duration`
cap and a
441 **sustained-inactivity split** to `trajectory_to_action_windows` ? port See et al.'s
rule (split /

```

```

442     end a window when ~80% of consecutive frames over a trailing window show no shuttle
motion).
443     Config-gated, default preserves current behavior. **Success:** segmental Edit + F1@k
improve vs
444     `CONTROL` on the golden set (?F1 CI excludes 0 **and** sign-test agrees) with **no
served-windowing
445     regression**; merge-rate drops materially. *(Win ? §5.2 See et al., ~95% end-of-rally
acc.)*
446     - **W2 · Rally-state features into the learned scorer (no special-casing). ** Wire the
already-built,
447     default-OFF cues as scorer feature columns: **net-crossings** (Gate-A), **two-sided
motion**, and a
448     **new cheap trajectory feature ? direction-reversal / peak count** (Hsu) as a hit-
cadence proxy;
449     audio hit-rate when the hit detector lands. **Re-measure with W1 in place and on the
served
450     windowing** ? recall person was ?0.021 and audio +0.079 but **both CIs spanned 0* at
n=6 on the
451     **pre-fix over-merged* candidates; the picture may change once windows are bounded.
**Success:** the
452     ablation (`ablation.py`) shows a cue with ?F1 CI clearing 0 + sign-test, surviving
served re-check
453     ? promote that cue; others stay default-OFF + retained (no deletion on a small-n
null). *(Win/medium.)*
454     - **W3 · Keep the scorer parameter-light (a discipline, documented not coded). ** Per
TemporalMaxer's
455     ablation (learned temporal conv overfits; parameter-free wins), resist replacing the
numpy
456     LogisticRegression with a learned temporal net until the corpus is much larger. Record
this as a
457     standing decision. *(Win ? §5.1 TemporalMaxer.)*
458
459     ### Tier 2 ? Medium effort
460     - **M1 · Serve-start classifier ? rally-start boundary** (See et al. pose/court, 88.6%).
Needs the
461     person cue + a **court-mask source** (manual polygon vs auto-homography ? **owner-
gated* per
462     [MULTI_SIGNAL_FUSION_PLAN.md])(MULTI_SIGNAL_FUSION_PLAN.md) §9.1). **Success:**
start-boundary
463     timing error drops on golden; promotes only if it lifts segmental F1@k without served
regression.
464     - **M2 · Per-venue threshold calibration** against golden labels (the heuristic's
footage-fragility ?
465     it scored 0.157 on multi-court). Tune `velocity_threshold` + `merge_gap` so local
segment counts
466     track golden counts per venue. **Success:** multi-court F1 rises toward the Gen-0
head's 0.561
467     without harming clean-court F1. *(Velocity is already fps-normalized ? calibrate per-
venue, not
468     per-fps.)*
469
470     ### Tier 3 ? Research bets (durable; corpus-gated)
471     - **R1 · Owned segmenter with an explicit boundary branch** (the ASRF/TriDet **idea* ? a
boundary head
472     that predicts cut points ? ported into our owned TAL head, NOT the full deep model,
and **trained**,
473     not grafted (§5.5 refutation)). Per
[OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md):
474     **TAL head first, VLM later.** **Gate:** the central unverified assumption is whether
boundary-head
475     gains hold at our scale on trajectory features ? so **pilot only when the corpus is
materially
476     larger (~?16?30 matches)**, and only after W1/W2 establish the heuristic ceiling.
**Success:**
477     beats the tuned heuristic + LogReg on held-out matches (segmental Edit + F1@k),

```

honestly LOVO.

```
478 - **R2 · Timestamp/sparse supervision** to get more labeled rallies per annotator-hour
($5.3) ? pilot
479 if/when we train a frame-wise model; research bet (residual gap, unproven on
trajectories).
480
481 ### Throughout ? corpus growth + active learning + label quality
482 - **Grow the golden set 6 ? ~16** via Roora ([`GOLDEN_SET_VENDORING.md`](archives/roora-
vendor/GOLDEN_SET_VENDORING.md)) ?
483 the single highest-leverage lever ([`NEXT_STEPS.md`](NEXT_STEPS.md)).
484 - **Active learning:** select the next videos to label by max **local-vs-Gemini
disagreement** ($4.4)
485 + scorer uncertainty + venue/fps diversity (our pragmatic stance; $5.3 flags the
formal methods as
486 an open research pass).
487 - **Measure IAA / label noise** as the corpus grows (Roora A/A overlap); revisit over-
segmentation-
488 robust losses only when a frame-wise model is on the table ($5.4 open question).
489
490 ### Dead-ends / do-NOT (so we don't relitigate)
491 - Don't treat **causal attention** (CausalTAD) as the over-merge fix ? it has no
boundary-split ($5.1).
492 - Don't expect **ASRF's boundary branch as a post-hoc drop-in** on our dense signal ?
refuted ($5.5).
493 - Don't adopt a **learned temporal net** at n=6?16 ? TemporalMaxer evidence says it
overfits ($5.1).
494 - Don't chase **bare tIoU-F1 / frame accuracy** ? blind to the over-merge ($5.4).
495 - Don't let **Gemini outputs become training data** for owned weights ? inference/shadow
only ($8).
496
497 ---
498
499 ## 7. Project: plan, tracking & definition of done
500
501 This roadmap is run as **one tracked project ? "Bridge the local?Gemini rally-detection
gap."** It
502 **completes when all committed tiers land AND a single quality number is attributed to
the effort**,
503 measured honestly and logged. This $7 is the living burn-down (update the status column
as PRs merge;
504 log each result in [`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md), numbers in
505 `eval_baselines/experiment_leaderboard.json`).
506
507 ### 7.1 The number (what we attribute to this effort)
508 Measured **leave-one-video-out, grouped by match, on the golden set** (grown to ?~16):
509 - **Headline:** multi-court rally-segmentation **tIoU-F1**, `baseline ? final`, plus
**gap-closure %**
510 = `(final ? baseline) / (Gemini ? baseline)`. Today: heuristic **0.157** ? Gemini
**0.66**
511 (Gen-0 head already **0.561**). **Success bar (proposed, owner-adjustable): close ?
50% of the
512 heuristic?Gemini multi-court gap** ? i.e. local multi-court F1 ? ~0.41, ideally
beating Gen-0's 0.561.
513 - **Primary lens (over-segmentation-sensitive, the actual failure):** segmental
**F1@0.25** +
514 **Edit-score** + **merge-rate** ? these must improve even if tIoU-F1 moves little.
515 - **Reported with honest stats** (bootstrap CI + paired sign-test); a result counts only
if it also
516 holds on the **served** windowing (eval?serve gate). The **baseline is frozen at E1**
before any fix.
517 - **Tracked headline = the n=6 golden *aggregate* tIoU-F1@0.5** (LOVO; the robust unit,
not a single
518 video) via `python -m backend.eval.rally_seg_eval`. The 0.157/0.66 figures above are
per-venue
519 *references*; computing an exact "% gap-to-Gemini closed" needs Gemini measured on the
```

```

*same* golden
520     set + metric ? **TODO** (don't mix the per-venue refs with the aggregate).
521
522     **Progress (cumulative, golden n=6 aggregate tIoU-F1@0.5):**
523
524     | after | F1 | merge_rate | ? vs baseline |
525     |---|---|---|---|
526     | E1 baseline (current heuristic) | 0.300 | 0.523 | ? |
527     | + W1 bounded windowing (cap=12, **default-OFF**) | **0.439** | **0.038** | **+0.139**
(6/0 sign-test) |
528     | + W2 net-crossings gate (mc=1, **default-OFF**) | **0.457** | 0.038 | **+0.157
cumulative** (W2 step +0.019, 6/0) |
529     | + W1.5 hard-cap fallback (**default-OFF**) | 0.444 | 0.000 | **+0.006 (n.s., CI spans
0)** ? neutral on golden; real-footage fix (mahadevpura-1 174s blob ? 24x?9s windows) |
530     | + W2 multi-feature scorer (person/audio) | _next (GPU)_ | | |
531
532     ### 7.2 Work items (status checklist)
533     Status: ? not started · ? in progress · ? done. (Each row = its own default-OFF,
measured PR.)
534
535     | ID | Deliverable | Tier | Depends on | Acceptance / promote criterion | Status |
536     |---|---|---|---|---|---|
537     | **E1** | merge/split-rate metric + F1@k surfaced + one-command `rally_seg_eval`
entrypoint; **froze baseline** (#185) | 0 | ? | over-merge shows up as a number on the 6 golden
videos | ? baseline F1 0.300, merge_rate 0.523 |
538     | **W1** | Bounded windowing ? `max_window_duration` cap + inactivity split (See et al.)
(#187) | 1 | E1 | seg F1@k ? vs baseline (CI?0 **and** sign-test), no served regression; merge-
rate ? | ? **?F1 +0.139 (6/0), merge_rate 0.523?0.038** (default-OFF) |
539     | **W1.5** | Hard-cap fallback (`window_hard_cap`) ? chop continuously-active windows at
the cap when no inactivity gap exists | 1 | W1 | golden ?F1 ?0 (or neutral) **** fixes a real-
footage over-merge it can't otherwise cap | ? **golden ?F1 +0.006 (n.s.)**; real `mahadevpura-1`
174s blob ? 24x?9s; retained **default-OFF** (real-footage + true-cap rationale, not a golden
lift) |
540     | **W2** | Rally-state features (net-crossings ? #188 + **never-empty safeguard** ?;
two-sided motion, **direction-reversal/peak**, audio = follow-on) | 1 | E1, W1 | ?1 cue ?F1 CI
clears 0 + sign-test, survives served re-check ? promote; others retained default-OFF | ? **net-
crossings ?F1 +0.019 (6/0)**; gate now **do-no-harm** (never zeroes a video); multi-feature
scorer next (GPU) |
541     | **W3** | Keep scorer parameter-light (standing decision; TemporalMaxer evidence) | 1 |
? | documented; no learned temporal net until corpus ? now | ? |
542     | **M1** | Serve-start classifier ? rally-start boundary (pose/court) | 2 | W1, **court-
mask decision (owner-gated)** | start-boundary timing error ?; promote on seg-F1 lift, no served
regression | ? |
543     | **M2** | Per-venue threshold calibration vs golden | 2 | E1 | multi-court F1 ? toward
Gen-0 0.561 without harming clean-court | ? |
544     | **R1** | Owned segmenter with an explicit **boundary head** (port ASRF/TriDet *idea*,
trained ? not grafted) | 3 (stretch) | corpus ?~16?30, W1/W2 | beats tuned heuristic+LogReg on
held-out matches (seg Edit + F1@k), LOVO | ? |
545     | **R2** | Timestamp/sparse-supervision pilot (label-efficiency) | 3 (stretch) | R1 path
| measured label-efficiency win | ? |
546     | **C** | Golden corpus 6 ? ~16 (Roora) + active-learning selection + measure IAA |
ongoing | ? | corpus ?16; IAA recorded | ? (~6 now; ~10 by the weekend) |
547
548     ### 7.3 Sequencing
549     `E1` (metrics + baseline) ? `W1`/`W2`/`W3` (the high-confidence wins) ? `M1`/`M2` ?
`R1`/`R2`
550     (corpus-gated). `C` (corpus growth) runs in parallel throughout ? it is the single
highest-leverage
551     input and gates the Tier-3 bets.
552
553     ### 7.4 Definition of done
554     The project is **complete** when:
555     1. **Tier 0 + Tier 1 + Tier 2** work items are merged ? each a measured PR (default-OFF,
promoted only
556     on measured lift, re-confirmed on the served windowing).

```

```

557     2. **The headline number (§7.1) is recorded** in `QUALITY_ITERATIONS.md` + the
leaderboard, with honest
558     stats, on the grown golden set ? **and the gap-closure target is met, or the honest
null is recorded
559     with a reasoned go/no-go on Tier 3.**
560     3. Each landed item left the harness greener (no regression) and `QUALITY_ITERATIONS.md`
tells the story.
561
562     > **Tier-3 scope decision (owner-adjustable):** R1/R2 (the owned boundary-head model +
sparse
563     > supervision) are **corpus-gated stretch goals** ? the verified research says deep TAL
detectors risk
564     > overfitting at n=6?16. **Default:** the project *completes* at Tier 0?2 + an
attributed number, and
565     > Tier 3 spins out into the owned-model milestone
([`OWNED_MODEL_IMPLEMENTATION_PLAN.md`](OWNED_MODEL_IMPLEMENTATION_PLAN.md)).
566     > Fold Tier 3 into the DoD instead if you want the project to span the owned model.
567
568     ---
569
570     ## 8. Constraints & guardrails (non-negotiable) + cross-links
571
572     - **Licensing:** every dependency ? code, data, AND weights ? must be **MIT / Apache-2.0
/ BSD**.
573     Reject viral/NC/custom (Gemma-3, Llama-4 MAU cap, Commons-Clause).
574     - **Never train owned weights on paid-LLM (Gemini/OpenAI/Anthropic) outputs**
(terms/copyright);
575     paid LLMs are inference/serving/shadow-reference only. Also forbidden: public datasets
known to
576     be LLM-generated. **Exclude minors** from any training pool; per-user data needs
explicit opt-in.
577     - **Owned model:** TAL head FIRST (de-risked, trains in minutes, doubles as a pseudo-
label teacher +
578     honest baseline the VLM must beat), then Qwen3-VL (Apache-2.0) / InternVL3 (MIT) ? see
579     [`OWNED_MODEL_IMPLEMENTATION_PLAN.md`](OWNED_MODEL_IMPLEMENTATION_PLAN.md).
580     - **Cross-links:** [`HOW_RALLY_DETECTION_WORKS.md`](HOW_RALLY_DETECTION_WORKS.md) .
581     [`EXPERIMENT_HARNESS.md`](EXPERIMENT_HARNESS.md) .
[`MULTI_SIGNAL_FUSION_PLAN.md`](MULTI_SIGNAL_FUSION_PLAN.md) .
582     [`ANALYZERS_AND_RECONCILERS.md`](ANALYZERS_AND_RECONCILERS.md) .
[`ABLATION_LAYER.md`](ABLATION_LAYER.md) .
583     [`EVALUATION.md`](EVALUATION.md) . [`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md) .
584     [`GOLDEN_SET_VENDORING.md`](archives/roora-vendor/GOLDEN_SET_VENDORING.md) .
[`OWNED_MODEL_IMPLEMENTATION_PLAN.md`](OWNED_MODEL_IMPLEMENTATION_PLAN.md).
585

```

8. assistant (2026-06-21T13:56:27.855Z)

I have the core doc. The §7 tracker is rich. Let me read the remaining track docs to confirm status and capture details, especially R2/R4, Gate-A wiring, and the experiment harness.

9. user (2026-06-21T13:56:28.761Z)

```

1     # R2 ? Task-faithful rally-boundary eval metric (tolerance-match + over-seg guard)
2
3     **Status:** **IMPLEMENTED (augmenting), 2026-06-17** ? owner-reviewed + decisions locked
2026-06-17;
4     `backend/eval/tolerance_metrics.py` + `rally_seg_eval` reporting landed (default-
additive; tIoU stays the
5     gate pending the ablation experiment). Gates the rally-end work (R4) and the dual-eval-
set regression gate.
6     A focused increment of the rally-quality project
([`RALLY_QUALITY_RESEARCH.md`](RALLY_QUALITY_RESEARCH.md)) §7;
7     the day-1 validation that motivates it:
[`REAL_FOOTAGE_VALIDATION_REPORT.md`](REAL_FOOTAGE_VALIDATION_REPORT.md);
8     the directive: `eval-boundary-tolerance` + `eval-dual-set-regression`.

```

9

10 > **Implementation refinement (2026-06-17):** §2.3.2's boundary-error distribution is computed over

11 > **unconditional best-overlap** matches, NOT the within-? 1:1 matches ? a ?-gated set is survivorship-biased

12 > (loose ends fail the gate and vanish, hiding the very end-deficit the diagnostic exists to surface; that

13 > within-? deficit instead shows up as low `tol_recall`). Matches the day-1 validation method. Measured on the

14 > n=9 corpus: `**|?start| 1.25 s vs |?end| 2.19 s**` ? "start?solved, end=the gap", confirmed.

15

16 ---

17

18 ## 0. TL;DR

19 `tIoU-F1@0.5`` is the wrong ruler for rally detection, in **both** directions: it over-penalizes the

20 **intrinsically fuzzy rally START** (the ~1?1.5 s service window), and its overlap matching is murky about

21 **over-production**. R2 replaces the headline with a **tolerance-match** metric under strict 1:1 (Hungarian)

22 assignment**, reported as a **decomposition** ? `P@? / R@? / F1@?` + **boundary-error** distributions split

23 into start/end** + an **over-segmentation guard** + a **??-sweep** ? not a single number.

24

25 **Honest**, slightly uncomfortable consequence (measured, see §3):** under R2 the local detector's headline

26 number *drops* and the gap to Gemini *widens* vs what tIoU showed ? because R2 faithfully charges for the

27 1.5× over-production and the loose ends that tIoU's overlap matching partly absorbed. That is the point:

28 R2 splits "behind Gemini" into its two real causes ? **(a) over-production [precision]**** and **(b) loose**

29 rally-ends [the R4 lever]** ? and forgives only the inherent start fuzziness. It is a more honest ruler,

30 not a flattering one.

31

32 ---

33

34 ## 1. Why `tIoU@0.5`` is the wrong ruler (grounded in the n=2 ground-truth clips)

35 1. **Over-penalizes the START.** For a 5.7 s rally, `tIoU ? 0.5`` tolerates only ~±1.9 s of boundary shift

36 (`? = D·(1??)/(1+?)``). The service window (stance ? racquet contact) alone is ~1?1.5 s ? even **Gemini's**

37 median `|?start|` is 0.98?1.46 s. So tIoU spends most of its budget on a boundary that is **intrinsically**

38 undefined to ~1 s. (Owner directive: don't strictly align starts; tolerate ~±2 s.)

39 2. **Count parity is a deceptive proxy.** mahadevpura-2 hit 1.03× of Gemini's count yet `F1@0.5 = 0.150` ?

40 right count, wrong boundaries. A count- or overlap-based view hides this.

41 3. **It flips config rankings on noise.** The hard-cap was golden-neutral at n=6 (+0.006 n.s.) but "best" at

42 n=7 once a continuously-active clip entered ? a metric artifact at the 0.5 cliff, not a quality change.

43 4. **The deficits we must see are start?end-symmetric.** Local's `|?start|` ? Gemini's; the **entire** gap is

44 the rally-**END** (`|?end| 2.1?3.1 s` vs Gemini ~0.8?1.2 s). A single blended IoU cannot show that ? and

45 directing R4 needs exactly that decomposition.

46

47 ---

48

49 ## 2. The metric

50

51 ### 2.1 Match rule (asymmetric tolerance ? forgive the start, expose the end)

52 A predicted window matches a golden rally iff `***|?start| ? ?_start` AND `|?end| ? ?_end`***`.

53 - `***?_start` generous (~2.0 s)**` ? absorbs the service-window/annotator start fuzziness (the part that is

54 `*not* a model deficit`).

55 - `***?_end` tighter (~1.5 s, but IAA-calibrated ? see §2.4)**` ? the rally-end IS a sharp physical event

56 (shuttle down/net/out), so a loose end `*is*` a real deficit we want the metric to expose, not forgive.

57

58 > A `*symmetric* ±2 s tolerance` is tempting (the owner's stated bar) but the preview (§3) shows it is `***not***`

59 > automatically more lenient than `tIoU 0.5` and it blurs exactly the end-deficit we need to see. Asymmetry is

60 > the point: relax the dimension that's intrinsically fuzzy (start), keep honest the one that's a real event (end).

61

62 `### 2.2 Assignment ? strict 1:1 (Hungarian), never overlap/greedy`

63 Match references to predictions by `***global Hungarian assignment***` (max total matches under the tolerance

64 gate), one-to-one. This is what makes `***over-production cost precision directly***`: ``P@? = matched / n_pred``

65 crashes when the detector emits 75 windows for 52 rallies. (tIoU's overlap matching let extras "match

66 something"; the W1 over-segmentation pendulum was barely costed.)

67

68 `### 2.3 Report the DECOMPOSITION, never one number`

69 1. `***P@?`, `R@?`, `F1@?`***` ? precision (over-production), recall (did we find the rally within tolerance), F1.

70 2. `***Boundary-error distributions***` ? median `***` IQR`***` of signed `*and*` absolute ``?start`, `?end`` over matched

71 pairs, `***split***`. This is the metric that says `***"start solved, end open"***` and aims R4.

72 3. `***Over-seg guard (mandatory):***` ``split_count`` (preds covering one ref ? W1 over-seg) and ``merge_count``

73 (refs under one pred ? the baseline mega-window). The `*original*` rule was strict per-video: `***a promotion`

74 may not increase either on ANY video.`***` `***REFINED` to a NET guard, 2026-06-17 (R1, owner-approved)`***` ? the

75 strict rule is correct for an `*incremental*` cue but `***mis-fires` on a structural mega-window?bounded change`***`

76 (splitting a mega-window necessarily lifts ``split_count`` from ~0, and the raw merge `*count*` is non-monotonic:

77 one window hiding 40 rallies is ``count=1, rate=1.0``; bounding it into pieces that each glue 2 neighbours is

78 ``count=9, rate=0.5`` ? FEWER rallies hidden). The net guard scores over-seg on the `***normalized rates***` as one

79 weighted cost per video (``w_merge.merge_rate + w_split.split_rate``, merges weighted higher ? a merge hides a

80 rally = recall loss, a split only fragments one), and `***passes` iff (a) the corpus-mean net cost does not rise

81 AND (b) no video's ``merge_rate`` regresses.`***` Implemented as ``tolerance_metrics.over_seg_guard`` (it always

82 reports the strict verdict too; revert = read ``strict_passes``). See QUALITY_ITERATIONS.md "R1" for the measured

83 verdict (milestone: mean ``merge_rate`` 0.694±0.150, net PASS, strict BLOCK).

84 4. `***?-sweep***` (``? ? {1.5, 2.0, 2.5, 3.0}``, symmetric for the trend; asymmetric pair for the headline) ?

85 because the convention/IAA noise is large, a single ? misleads (§3). Report the curve.

86 5. `***tIoU`/mIoU`` demoted to a SECONDARY trend-line diagnostic`***` (keeps the paper-ablation backbone

87 comparable) ? `***not the gate.***`

88

89 `### 2.4 Choosing ? ? it MUST be calibrated to inter-annotator agreement`

```

90     The preview (§3) shows even Gemini scores `tolF1@1.5 = 0.20` on mp2 ? i.e. on ~half
the rallies the
91     golden END differs from Gemini's by >1.5 s. That is plausibly label/convention noise,
not model error**:
92     we'd be measuring how two humans (or human-vs-Gemini) disagree on "when did the rally
end." Therefore
93     `?_end` must be ? the inter-annotator end-agreement**, which we do not yet know.
Action: a small IAA
94     study ? a 2nd labeler re-labels 2?3 already-golden videos; set `?_start`/`?_end` to
~the 90th-percentile
95     pairwise human |?start|/?end|. Until then freeze provisional ?_start=2.0, ?_end=1.5
and always report the
96     sweep**, with the `? = D.(1??)/(1+?)` derivation recorded so the choice is auditable.
97
98     ---
99
100    ## 3. Preview on the two ground-truth clips (greedy-1:1 approximation; Hungarian in the
impl)
101
102    Symmetric-? sweep, F1@? (local-best = the per-clip best config: mp2 W1+W2+hardcap,
GX010128 W1+W2):
103
104    | clip / method | tolF1@1.5 | @2.0 | @2.5 | @3.0 | (tIoU-F1@0.5) |
105    |---|---|---|---|---|---|
106    | mahadevpura-2 / Gemini | 0.203 | 0.354 | 0.506 | 0.557 | 0.557 |
107    | mahadevpura-2 / local | 0.111 | 0.148 | 0.222 | 0.278 | 0.296 |
108    | GX010128 / Gemini | 0.615 | 0.769 | 0.813 | 0.857 | 0.813 |
109    | GX010128 / local | 0.189 | 0.315 | 0.346 | 0.362 | 0.457 |
110
111    Reading the preview (all honest):
112    - tolF1@~3.0 ? tIoU-F1@0.5 ? the sweep is a smooth, interpretable knob that
contains tIoU as a point;
113    good (continuity with the frozen baseline).
114    - The local?Gemini gap WIDENS under the tighter, 1:1 tolerance (GX010128 ratio 0.56
at tIoU0.5 ?
115    0.41 at ?2.0). R2 charges local for the 1.5× over-production (precision) and the
loose ends ? exactly
116    the deficits tIoU under-counted. This is the metric working, not local regressing.
117    - Even Gemini drops at small ? (mp2 0.557?0.203 from ?3.0?1.5) ? strong evidence the
golden END
118    convention carries >1.5 s of noise ? §2.4's IAA calibration is a prerequisite, not a
nicety.
119    - Boundary-error medians (the robust diagnostic, from the validation report): local
|?start| 0.78?1.74 s
120    (? Gemini 0.98?1.46) vs local |?end| 2.12?3.12 s (? Gemini 0.76?1.19) ? the headline
R4 signal.
121
122    ---
123
124    ## 4. Promotion / "trustworthiness" rule (small-n)
125    A candidate is promotable only if all hold:
126    1. Paired per-video sign-test sweep on F1@? is near-unanimous: 6/0 @ n=6
(p=0.031) or 7/0 @ n=7
127    (p=0.016) (8/0 @ n=8, p=0.008); 5/1 or 6/1 is NOT significant. Run on the
growing-golden set.
128    2. No over-seg regression ? the net over-seg guard passes (§2.3.3): the corpus-
mean weighted
129    merge/split-rate cost does not rise AND no video's `merge_rate` regresses. (Was the
strict per-video
130    count rule; refined to net on 2026-06-17 ? see §2.3.3.)
131    3. Effect size + bootstrap CI reported as an honesty band (expect it to span 0 at
this n; use it to
132    separate "real" from "directionally-clean-but-negligible," not as the pass/fail).
133    4. Stratify, don't blend ? report the continuously-active blob clip (mahadevpura-1)
as its own stratum

```



```

134     (its near-degenerate scores can swing a small-n mean).
135     5. **No regression on EITHER eval set** (raw + growing-golden) ? the dual-set
discipline.
136
137     ---
138
139     ## 5. Implementation plan (pure, offline, additive ? zero production risk)
140     - **`backend/eval/tolerance_metrics.py`** (pure, GPU-free, unit-testable):
141     - `match_1to1(preds, gts, ?_start, ?_end) -> (matches, P, R, F1)` ? Hungarian via
142     `scipy.optimize.linear_sum_assignment` if available, else a pure  $O(n^3)$  fallback
(corpus n is tiny).
143     - `boundary_error_report(preds, gts) -> {?start, ?end: signed/abs median+IQR}`
(computed over
144     unconditional best-overlap matches internally ? see the refinement note above).
145     - `over_seg_report(preds, gts) -> {split_count, merge_count}` (reuse the bipartite
logic from
146     `segmentation_metrics.merge_split_report`).
147     - **Wire into `backend/eval/rally_seg_eval.py`** ? emit the R2 block **alongside** tIoU
(don't remove tIoU);
148     add `--tau-start/--tau-end` (both landed; the ?-sweep helper
`tolerance_metrics.tolerance_sweep` exists
149     but is not yet surfaced as a `rally_seg_eval` CLI flag). Surface in the experiment
leaderboard.
150     - **`run_regression.py`** ? add the dual-set R2 aggregate + the over-seg-guard gate.
*(Status: the guard
151     primitive landed as `tolerance_metrics.over_seg_guard` and is wired into
`rally_seg_eval.compare()`/`--vs`
152     ? the natural home, since `rally_seg_eval` loads the golden manifest + trajectories.
The dual-set
153     `run_regression.py` aggregate is still a separate future item.)*
154     - **Tests** ? synthetic intervals: exact-match, start-fuzzy-within-?, end-loose-
beyond-?, over-production
155     (precision crash), the Hungarian-vs-greedy tie case, empty inputs.
156     - **Default-additive ? ablation-gated replacement:** R2 is reported next to tIoU
(augment). The *gate*
157     switches to R2 only after a dedicated **ablation experiment** on the golden corpus
shows R2 ranks candidates
158     at least as sensibly as tIoU ? so the promotion bar never changes silently mid-flight.
When tIoU is retired
159     as the gate it is **deprecated with documentation, not deleted** (kept as the
secondary trend-line + the
160     `QUALITY_ITERATIONS.md` paper-ablation backbone, with a dated note recording the
ablation evidence for the
161     switch). Owner decision, 2026-06-17 (§7).
162
163     ---
164
165     ## 6. Definition of done
166     - `tolerance_metrics.py` + tests merged (default-additive); `rally_seg_eval` reports the
R2 decomposition
167     next to tIoU (the ?-sweep helper exists in `tolerance_metrics` but is not yet wired
into the
168     `rally_seg_eval` report); the corpus (n?8) re-scored under R2 and recorded in
QUALITY_ITERATIONS.
169     - The IAA study (§2.4) scoped (it can land after R2's code; R2 ships with provisional ?
+ the sweep).
170     - `run_regression.py` carries the dual-set R2 gate + the per-video over-seg guard ? but
R2 **augments** tIoU
171     until the ablation experiment (§7) validates it as the gate; only then does the gate
switch.
172     - tIoU/mIoU retained as a secondary trend-line; when retired as the gate, **deprecated-
with-docs (kept for
173     posterity)** per the §7 owner decision ? a dated note records why and links the
ablation evidence.
174

```

```

175    ---
176
177    ## 7. Decisions
178    **Locked (recommended, design-level):** Hungarian 1:1 (not greedy/overlap); report the
decomposition + sweep,
179    never one number; tIoU demoted to secondary; R2 default-additive until validated;
boundary-error split
180    start/end is the R4-aiming diagnostic.
181
182    **Locked (owner review, 2026-06-17):**
183    - **?_end asymmetric (1.5 s) vs ?_start (2.0 s) ? ADOPTED.** Asymmetry is the point:
forgive the intrinsically
184    fuzzy start, keep the rally-end honest. Your "±2 s, not more" bar maps to
`?_start=2.0`; `?_end` stays the
185    tighter, IAA-calibrated knob (§2.4).
186    - **IAA study after R2 code ? ADOPTED.** R2 ships now with provisional `?_start=2.0 /
?_end=1.5` + the full
187    sweep; the IAA study lands next to fix ?. R2's code is not blocked on it.
188    - **Gate: AUGMENT now, then ablation-gate the eventual replacement, then deprecate-with-
docs ? ADOPTED.**
189    R2 is reported **alongside** tIoU (augment) until validated. tIoU is **not** swapped
out silently: before
190    R2 becomes the promotion gate, run a dedicated **ablation experiment** on the golden
corpus showing R2
191    ranks candidates at least as sensibly as tIoU (and exposes the deficits tIoU hid).
When tIoU is finally
192    retired as the gate, it is **formally deprecated with documentation ? kept for
posterity, never deleted**:
193    it remains a secondary trend-line and the paper-ablation backbone
(`QUALITY_ITERATIONS.md`), with a dated
194    note recording *why* it was superseded and the ablation evidence behind the switch.
(Consistent with the
195    repo's "retain weak/superseded signals; numbers inform, foundations decide" norm.)
196

```

10. user (2026-06-21T13:56:29.448Z)

```

1    # R1 Milestone ? Launch Report (windowing default flip: W1 cap=6 + R4)
2
3    **Status:** LAUNCHED 2026-06-18 (PR #204). **Flip:** the local-windowing defaults go ON
?
4    `max_window_duration 0?6`, `window_inactivity_end 0?1.5`, `inactivity_rally_thresh
0.08?0.20`,
5    `inactivity_min_density 0.34?0.30`. **W2 net-crossings, W1.5 hard-cap, and R5 blob-
fallback all stay OFF.**
6    This is the [Launch Report convention](../..) record made **before** the flip merged ?
total quality impact
7    on **both rulers** + the over-seg guard + honest caveats, for posterity.
8
9    > **Why this flip is safe to ship:** measured on the **n=13** golden corpus (the grown
set, incl. the two
10    > Boxhill clips GX010137/GX010141), the milestone clears the **net over-seg promotion
bar** with room to
11    > spare, and the verdict was **independently reproduced** (a different harness, #214)
and **adversarially
12    > stress-tested** (5-lens L00/seed/per-clip workflow). Nothing in the verification
refuted it.
13
14    ---
15
16    ## 1. Total quality impact ? BOTH rulers (n=13, leave-one-video-out unit = whole video)
17
18    The headline promotion metric is the **R2 task-faithful tolerance-match F1** (`tolF1`,
asymmetric ?_start=2/
19    ?_end=1.5); **tIoU-F1@0.5** is kept as the secondary continuity ruler.

```

```

20
21 | config | **R2 tolF1** | **tIoU-F1@0.5** | mean \[?end\] |
22 |---|---|---|---|
23 | baseline (all-OFF) | 0.137 | 0.236 | 44.6 s |
24 | + W1 cap=12 | 0.210 | 0.382 | 5.7 s |
25 | + cap=6 (R3) | 0.320 | 0.436 | 5.0 s |
26 | **+ R4 (milestone, shipped)** | **0.359** | **0.474** | **3.3 s** |
27
28 **Paired sign-tests (the promotion evidence):**
29
30 | transition | ? tolF1 (95% CI) | sign (W/L/T) | p | ? tIoU-F1 | sign | p |
31 |---|---|---|---|---|---|---|
32 | **baseline ? milestone** | **+0.222** [+0.152, +0.286] | **12/0/1** | **0.0005** |
+0.238 [+0.159, +0.318] | 12/0/1 | 0.0005 |
33 | cap12 ? cap6 (R3) | +0.110 [+0.070, +0.153] | 12/0/1 | 0.0005 | ? | ? | ? |
34 | cap6 ? +R4 | +0.040 [+0.015, +0.065] | 9/1/3 | 0.0215 | ? | ? | ? |
35
36 The single non-win in baseline?milestone is **mahadevpura_1** (the continuously-active
blob: tolF1 0.000?0.000,
37 a tie ? it needs R5, which stays OFF). **Both new clips improve:** GX010137 0.448?0.685,
GX010141 0.333?0.431.
38
39 ## 2. Over-segmentation guard (the promotion gate)
40
41 The R1 **net** over-seg guard (`tolerance_metrics.over_seg_guard`) is the gate (not the
strict per-video count
42 rule). Baseline ? milestone:
43
44 - **NET verdict: PASS** ? weighted (merge:split = 2:1) merge_rate/split_rate cost
**1.303 ? 0.429** (? ?0.874).
45 - **mean merge_rate 0.651 ? 0.161** ? drops on **all 13** videos; **zero per-video
merge_rate regressions**.
46 - **strict_passes = False BY DESIGN** (not a regression): the raw per-video *count* rule
trips because splitting
47 a mega-window lifts split_count from ~0 and the raw merge *count* is non-monotonic
(one window hiding 40
48 rallies = count 1/rate 1.0 ? bounded into pieces gluing 2 neighbours = count 9/rate
0.5 ? **fewer** rallies
49 hidden). The guard scores the recall-meaningful *rate*; see
[QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) "R1".
50
51 ## 3. Verification provenance (independent + adversarial)
52
53 - **Independent reproduction (#214 `lovo_resweep.py`, a different harness):** milestone
vs baseline 0.137?0.359,
54 sign **12-0-1, p<0.001**; R4 on-vs-off 0.320?0.359, sign **9-1-3, p=0.021** ? matches
to 3 decimals. Its LOVO
55 cap-selection picks **cap=6 unanimously (13/13)** on tIoU.
56 - **5-lens adversarial workflow** (reproduce / leave-one-out / R4-loss audit / new-clip
sanity / guard audit ?
57 synthesis): verdict **STANDS**. Every LOO fold (n=12) clears CI-excludes-0 AND p<0.05;
worst fold (drop
58 GX010128, the biggest gainer) still ?+0.202, p=0.0010 ? **no single clip drives the
win**.
59
60 ## 4. Honest caveats (recorded so they aren't relitigated)
61
62 1. **The W1 cap is the dominant lever; R4 is a marginal increment.** baseline?cap6 is
~+0.18 of the +0.222
63 tolF1 gain; the R4 inactivity-end trim adds a *small but significant* +0.040
(p=0.0215, sign 9/1/3). Frame R4
64 as "a small, significant end-precision gain" (mean [?end] 4.96?3.31 s), not a co-
equal driver.
65 2. **R4's one loss is GX010137** (tolF1 0.727?0.685, ?0.042) ? a tolerance-accounting
artifact, **not** an R4

```

```

66      failure: on that clip strict F1 *improves* (0.805?0.822), |?end| *improves*
(0.447?0.439), over-seg unchanged;
67      R4 simply trimmed 4 trailing predicted segments (43?39, toward gt=34). The clip's
*overall* milestone delta is
68      a large win (+0.237).
69      3. **GX010141 is the worst-behaved clip post-milestone** (merge_rate still 0.448, ~30 s
windows vs gt max ~10.5 s)
70      ? partially-merged, **not** a regression (it improves +0.098) and it does not break
the guard, but it's a
71      future over-merge target as the corpus grows.
72      4. **cap=4-vs-6 wobble on tolF1** (from #214): cap=4 edges cap=6 by ~0.005 (noise) ? but
**LOVO re-fitting to
73      cap=4 generalizes *worse*** (?0.004) and tIoU is firmly cap=6, so **cap=6 stands**.
Watch as n grows; switch
74      only if cap=4 pulls ahead *consistently across folds*.
75      5. **n=13 is small.** "Robust" here means "no single clip or bootstrap seed flips the
verdict," not large-sample
76      certainty. (Bootstrap CIs jitter ~0.001?0.004 run-to-run; means/sign-tests/p-values
are exact.)
77      6. **Nightly re-baselined + an all-OFF floor added (DONE).** The flip made the nightly's
`default` config
78      (= SERVED) equal `milestone`, so its config-drift detector flagged the frozen
baseline STALE ? re-snapshotting
79      in #216. A pinned **`baseline_off`** reference config was then added (follow-up) so
every nightly report keeps
80      showing the delta the milestone buys (floor tolF1 **0.137** vs milestone **0.359**) ?
the launched win has to
81      keep standing its ground, not be taken at launch-time value.
82      7. **Pure-logic eval only** (CPU, persisted trajectories) ? no GPU/real-video re-
detection was run; the flip
83      changes only the windowing knobs the served path already reads.
84
85      ## 5. What did NOT flip (still default-OFF, retained)
86
87      - **W2 net-crossings gate** ? no ground-truth-backed lift (the weak-signals-retained
directive keeps it).
88      - **W1.5 hard-cap** ? over-segments normal clips (chops before R4); superseded by R5 for
the blob case.
89      - **R5 blob-fallback** ? a stratum tool for continuously-active clips; turn on
selectively, not globally.
90

```

11. assistant (2026-06-21T13:56:32.825Z)

Now I have a clear picture of where R-series stands (R1 shipped, R2 implemented as augmenting). Let me quickly read the remaining track docs to confirm Gate-A wiring status, harness, golden fixtures, and the QUALITY_ITERATIONS for the latest measured state, plus check the \$7 tracker freshness against the docs.

12. user (2026-06-21T13:56:33.822Z)

```

1      # Multi-signal rally fusion ? add a human (person) cue alongside the shuttle
2
3      **Status:** **P1 DONE** (2026-06-13) ? the person cue is extracted to
4      `backend/pipeline/detectors/person_detector.py` (`PersonDetector`), consumed by
`yolo_hybrid` with no
5      behavior change. **P2?P4 remain owner-gated.** Phases below are sequenced but each is a
separate PR.
6      **Owner decisions captured** in [$9](#9-open-decisions-for-the-owner).
7
8      > **Companion docs:** the 2-layer mental model is
[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md);
9      > the modular-fusion architecture this builds on is **ADR-007**
([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md));
10     > the permissive-dependency rule is **ADR-002 / ADR-005** ; the audio cue (the #1 next

```

```

cue, already
11     > scaffolded) is [archives/research/AUDIO_RALLY_DETECTION_PLAN.md](archives/research/AUDIO_RALLY_DETECTION_PLAN.md).
12     > The vendor labels that gate the learned path come from
[GOLDEN_SET_VENDORING.md](archives/vendor/GOLDEN_SET_VENDORING.md)
13     > and the labeling spec in [ROORA_LABELING_GUIDELINES.md](archives/roora-vendor/ROORA_LABELING_GUIDELINES.md).
14     > **How each cue here is added/tested without production regression** ? A/B, shadow
eval, and the
15     > promote-only-on-lift discipline ? is [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)
(every P1?P4 cue
16     > below lands flag-gated, default OFF, promoted only on measured LOVO lift).
17
18     ---
19
20     ## 1. Context ? why this, why now
21
22     Rally detection today runs **one cue at a time**. Each segmenter independently turns a
single signal
23     into rally windows:
24
25     | Segmenter | Signal | File |
26     |---|---|---|
27     | `wasb_hybrid` / `tracknet_hybrid` | shuttle velocity (frozen WASB/TrackNet trajectory)
| `backend/pipeline/detectors/tracknet_runner.py` |
28     | `yolo_hybrid` (the Pydantic `default_segmenter` fallback; `config.json` overrides it
to `gemini`) | person velocity (torchvision + ByteTrack) |
`backend/pipeline/segmenters/yolo_hybrid.py` |
29     | `motion` | raw pixel motion | `backend/pipeline/segmenters/motion.py` |
30     | `gemini` | a cloud VLM watching the whole clip | `backend/providers/gemini.py` |
31
32     **There is no fusion.** The owner's directive: make the detector **use the shuttle *and*
the humans
33     together** ? and be ready to fold in more signals (audio, later pose) ? delivered
34     **MIT / Apache-2.0 / BSD-clean** so we can ramp commercially without a licensing
rethink.
35
36     This is **not greenfield** ? it operationalizes work the repo already blessed:
37
38     - **`HOW_RALLY_DETECTION_WORKS.md` open question:** **"Fusion: combine shuttle-motion
(WASB) with
39     player-motion for static cams."**
40     - **ADR-007** already designed a **modular "rally layer"** that fuses cues **outside* the
frozen WASB
41     model, with **audio as the #1 next cue and person/pose parked as #2**. This plan picks
up the parked
42     person cue and generalizes the fusion seam so future cues are "register a producer,"
not "re-touch
43     the segmenter."
44     - **`backend/pipeline/detectors/rally_gate.py` already exists** ? the net-crossing gate,
explicitly
45     labelled in its own docstring as **"Gate A in the over-fire fix spectrum (**B = player-
stance state
46     machine, future**)." The person cue in this plan **is that Gate B ? now built as the
`FusionSegmenter` person path (default-OFF + unmeasured on golden; see §10 P2).**
47
48     ### The owner's "held-shuttle" false positive ? what's actually going on
49
50     The owner observed a "rally" that was just a player **holding / flicking the shuttle in
hand** between
51     points. This is the **exact* failure `rally_gate.py` was written to kill:
52
53     > **"WASB over-detects: when a player flicks/tosses the shuttle back to the opponent for
service after
54     > losing a point, that single trajectory looks like a rally. The discriminator is

```

```

structural: a real
55     > rally has the shuttle crossing the net MULTIPLE times; a single service feed crosses
~once."*
56
57     **The structural fix already exists and is validated offline ? but it is NOT wired into
the production
58     runtime path.** `rally_gate.apply_gate(..., min_crossings=2)` is called only from the
**eval drivers**
59     (`eval_partial_labels.py`, `export_wasb_segments.py`, `calibrate_local.py`;
`distill_local.py` uses
60     `rally_gate.gate_windows` to compute the `net_crossings` feature, not `apply_gate`); the
live segmenters
61     historically didn't apply it. P2 closed this: `FusionConfig.min_crossings` now exists
(default `0` =
62     OFF), so the served Gate A is a config flag ? consistent with $10 P2 below, marked DONE.
So the
63     single cheapest, highest-confidence win in this whole plan is **P2's first step: fold
the existing
64     Gate A into production.** Everything else (the person cue / Gate B, the learned fusion)
is robustness
65     on top of that.
66
67     ### Why humans help ? and the honest risk
68
69     The beachhead videos are **static tripod** (Bangalore academies,
[PMF_MARKET_RESEARCH.md](archives/field-pmf/PMF_MARKET_RESEARCH.md)),
70     so player presence/motion is a usable cue. ADR-001's warning was specifically about
**broadcast pan**
71     cameras, where player-motion fooled the old heuristic ? *that failure mode does not
apply to a fixed
72     court cam.* But we keep **the shuttle as the camera-invariant primary** so the general
bet stays robust
73     on other camera types. Humans then add:
74
75     - **Precision** ? reject shuttle-on-floor / warm-up knock-ups / held-shuttle feeds when
the court
76     isn't occupied by ?2 players in a play posture.
77     - **Recall** (via the learned classifier) ? features that help bridge shuttle
occlusion/blur gaps.
78
79     ---
80
81     ## 2. Core architecture ? one rally-fusion layer, cues are swappable producers
82
83     Formalize ADR-007's diagram. Every cue is a **black box that emits a time-aligned
signal** (per-frame
84     activity and/or per-window features); **fusion happens in *our* layer, never inside a
frozen model.**
85
86     ```
87     frames ?? WASB (frozen, MIT)      ?? shuttle trajectory ???
88     frames ?? person detector (ours)  ?? tracks + boxes ??????
89     audio ?? hit detector (ours)      ?? hit events ???????????? RALLY LAYER (ours) ??
rallies
90     [frames ?? pose/stance (ours; parked #3) ?? stance ???????? feature + decision fusion
91     ```
92
93     - **Shuttle = primary, camera-invariant.** Windows still **seed** from
94     `TrackNetRunner.trajectory_to_action_windows()` (`tracknet_runner.py:234`). Other cues
*enrich and
95     adjudicate* shuttle-seeded windows; they do **not** replace the shuttle seed. Person
runs only over
96     **shuttle-seeded candidate windows + padding**, not the whole video ? which bounds the
added cost.
97     - **A small `RallyCue` contract.** Each producer yields `(frame_idx ? features)` + per-

```

```

window stats.
98     Adding a cue = register a producer; the existing `SegmenterRegistry`
(`segmenters/base.py:35`) is the
99     template for the pattern. Keep the single-cue segmenters registered as fallbacks.
100
101     ---
102
103     ## 3. The two fusion modes (both ? per the owner decision)
104
105     ### 3.1 Decision-level gate (safe baseline, ships first, no labels needed)
106
107     ```
108     rally_active = shuttle_motion ? (?2 persons in court zone ? recent person track)
109     ```
110
111     The `? recent-track` clause tolerates 1?2 frame person-detector misses so the AND-gate
    can't cause a
112     recall cliff. This composes with the existing net-crossing Gate A:
113
114     ```
115     keep window ? net_crossings ? 2                (Gate A ? EXISTS in rally_gate.py, eval-
    only today)
116             ? (?2 in-court players ? recent track)  (Gate B ? the person cue, BUILT
    default-OFF, unmeasured on golden)
117     ```
118
119     Gate A kills the service-feed / held-shuttle case structurally; Gate B adds court-
    occupancy evidence
120     for the cases a single trajectory can still fool (e.g. a knock-up that does cross the
    net). No training
121     data required ? this is the label-free fusion that ships before any golden labels exist.
122
123     ### 3.2 Feature-level (primary, the learned path)
124
125     The cues emit a small set of aggregated, scale/translation-invariant per-window
    features ? NOT raw
126     coordinates. Raw pixel coordinates would memorize this court/camera and fail to
    generalize off a
127     still-small labeled set; counts / ratios / zone-membership / two-sidedness transfer.
128
129     What already feeds the classifier (`distill_local.py:40`, `FEATURE_NAMES`):
130     `duration, peak_velocity, mean_velocity, active_ratio, net_crossings` ? note
    net_crossings is
131     already a feature** (computed via `rally_gate.gate_windows`).
132
133     Person features to add (~4, deliberately few):
134
135     | Feature | Meaning | Rally signal |
136     |---|---|---|
137     | `players_in_court` | median count of persistent in-court tracks | ? 2 (singles) / 4
    (doubles); 0?1 = dead time |
138     | `two_sided_motion` | are moving players on both net-halves? | rally = two-sided;
    one person feeding = one-sided |
139     | `mean_player_motion` | aggregate normalized player velocity | rally = sustained;
    resting/standing = low |
140     | `in_court_ratio` | fraction of frames with ?2 in-court players | track stability vs
    flicker |
141     | `shuttle_player_colocation` | fraction of frames the shuttle sits inside/adjacent to a
    person box (hand region) | held = high; in-flight = mostly free space ? the direct held-
    shuttle discriminator |
142
143     These five person features are now built
    (`fusion_features.py::PERSON_FEATURE_NAMES`; default-OFF, unmeasured on golden) and join the
    shuttle features in the candidate `stats` JSON
144     (`database.py::add_candidate_segment`) and the feature vector of the existing

```

```

distilled classifier ?
145     the tiny numpy-only L2 logistic regression in `backend/eval/classifier.py` (ADR-004),
NOT a new net.**
146     It scores each candidate window `P(real rally)` and replaces the Gemini call at runtime.
147
148     **What the model learns:** the **weights** on those few features ? e.g. down-weight a
shuttle-motion
149     window with 0?1 in-court players / one-sided motion / high shuttle-colocation (warm-up,
floor shuttle,
150     held-shuttle feed); up-weight ?2 players, two-sided, sustained motion, shuttle in free
flight.
151
152     **Small-N discipline (non-negotiable):**
153     - Keep to **~4 person features** ? more features on a still-small labeled set overfit;
that's exactly
154     why the model stays a logistic regression, not a net.
155     - Use the **leave-one-video-out (LOVO)** protocol that `distill_local.py` **already
implements**: train
156     on the other videos' candidates, predict the held-out *video*, score vs its labels ?
never a
157     held-out *window* from the same video (windows in one video are correlated; that
flatters the score).
158     - Training labels today are **Gemini silver labels** (distilled offline, ADR-004/007 ?
Gemini is a
159     teacher, never a runtime dependency, never a training-data *source* for owned
weights); **golden
160     labels are the honest measuring stick**, and as the Roora golden corpus grows they can
also become
161     training labels. Until the labeled set is larger, treat the learned weights as
162     **calibration / direction-check, not "a fully general model"*** ? the decision-level
gate (§3.1) is
163     the shippable, label-free fusion in the meantime.
164
165     > *(Gated on golden labels existing ? the same blocker ADR-007 flags for audio. See
166     > [GOLDEN_SET_VENDORING.md](archives/roora-vendor/GOLDEN_SET_VENDORING.md).)*
167
168     ---
169
170     ## 4. Shuttle-state: in-flight vs in-hand (the misclassification, mapped to real code)
171
172     The "shuttle moved" seed conflates a shuttle **in flight** with a shuttle **carried in
hand**. The
173     discriminators, and where each lives:
174
175     | Signal | In-flight | Held / static | Status |
176     |---|---|---|---|
177     | `net_crossings` | ? 1 (usually several) | ~0 (held) / ~1 (single feed) | **EXISTS**
(`rally_gate.py`); served Gate A wired via `FusionConfig.min_crossings` (default 0 = OFF) |
178     | `shuttle_player_colocation` | low (free space) | high (near a hand) | **BUILT**
(`fusion_features.py`; default-OFF under the person cue, unmeasured on golden) |
179     | `direction_reversals` | several (racket contacts) | ~0 | derivable from trajectory;
partial overlap with crossings |
180     | `peak shuttle speed` | high | walking speed | derivable from existing velocity |
181
182     So the in-flight-vs-held fix is **mostly Gate A (already built) + the new colocation
feature from the
183     person cue**. **No new label column is needed** ? the golden labels already start a
rally at the
184     *serve-contact frame* and treat all pre-serve / held-shuttle time as dead time (see
185     [ROORA_LABELING_GUIDELINES.md](archives/roora-vendor/ROORA_LABELING_GUIDELINES.md)), so
held-shuttle moments are negatives
186     by construction and are the ground truth that proves the detector wrong.
187
188     ---
189

```



```

190     ## 5. Adversarial mitigations (integrating safely)
191
192     - Court-region mask ? polygon around the court; drop persons outside it (spectators,
    coaches,
193     adjacent courts). Config-supplied per camera; ~1 ms. Source: a manual per-camera
    polygon first;
194     the 4 court corners + net-line that Roorla labels once per video (see the labeling
    doc) build
195     this mask and the net split that `two_sided_motion` / `net_crossings` need.
196     - Player-count constraint ? expect 2 (singles) / 4 (doubles) persistent in-court
    tracks;
197     suppress transient/extra detections. The `singles_or_doubles` manifest flag sets the
    expectation.
198     - Temporal association ? reuse `supervision` ByteTrack (already in
`yolo_hybrid`); reject
199     singleton detections; smooth across frames.
200     - Multi-court rejection ? proximity to the primary court region (homography
    optional, later).
201     - Domain shift (COCO frontal ? side/overhead court) ? start box-only (robust);
    keep the
202     detector swappable. Fine-tuning a person detector is parked ? we have no in-house
    person labels,
203     and the paid-LLM/training guardrail + minors-exclusion apply to any future person-
    training data.
204     - Fusion-degradation guard ? shuttle stays primary; person is gate + feature, never
    the sole
205     trigger; the single-cue segmenters remain registered fallbacks.
206
207     ---
208
209     ## 6. Box vs pose
210
211     Box-first. Occupancy + count + two-sided motion + shuttle-colocation is enough for
    both the gate
212     and the features, and box detection is the cheap, robust win. 2D pose is parked behind
    the same
213     `RallyCue` interface ? pose (MediaPipe Pose / RTMPose, both Apache-2.0) buys serve-
    ready stance
214     gating and shot cues later, at the cost of a second model. This matches ADR-007 parking
    pose as the
215     #2/#3 cue.
216
217     ---
218
219     ## 7. Performance
220
221     The person detector is the existing in-process torchvision model run at frame-skip
    375
222     (`yolo_frame_skip` default 3) with track interpolation between sampled frames. The
    shuttle keeps
223     running in its separate WSL process ? the two are separate processes, so a "shared
    backbone" isn't
224     free; budget realistically as the existing `yolo_hybrid` cost minus skip savings. Person
    runs only
225     over shuttle-seeded windows + `ai_handoff_padding`, not the whole video, which bounds
    the cost.
226
227     ---
228
229     ## 8. Licensing (the guardrail table ? all green)
230
231     Every dependency stays MIT / Apache-2.0 / BSD (ADR-002 / ADR-005; `ultralitics` was
    dropped for
232     AGPL-3.0).
233

```

```

234 | Cue / model | Code | Weights / data | Verdict |
235 |---|---|---|---|
236 | Shuttle: WASB-SBDT, TrackNetV3 | MIT | frozen badminton weights | ? (ADR-001/002) |
237 | Person (default): torchvision Faster R-CNN / RetinaNet / FCOS | BSD/MIT | COCO-
pretrained | ? already in repo (`_DETECTOR_FACTORIES`) |
238 | Person (escalation): **YOLOX**, **RT-DETR / RT-DETRv2** | Apache-2.0 | COCO /
Objects365 | ? |
239 | Pose (parked): **MediaPipe Pose**, **RTMPose / MMPose** | Apache-2.0 | COCO-structure
| ? |
240 | Tracking: `supervision` ByteTrack | MIT | n/a | ? already in repo |
241 | **Banned** | **Ultralytics YOLOv5/8/11 (AGPL-3.0)** . **MonoTrack (Adobe, non-comm)**
. **YOLO-NAS (Deci, non-comm)** | ? | ? |
242
243 **COCO caveat to record:** COCO *annotations* are **CC BY 4.0** (commercial OK *with
attribution*);
244 COCO *images* are Flickr-ToS (per-image). Pretrained-weight use is clean; record the
attribution.
245
246 ---
247
248 ## 9. Open decisions (for the owner)
249
250 1. Court-region mask source ? manual per-camera polygon (simple, ship first) vs auto
court
251 homography/line detection (later). Recommend manual first, fed by Roora's court-
corner labels.
252 2. Person detector default ? keep torchvision (zero new deps, BSD/MIT) vs adopt
YOLOX/RT-DETR
253 (Apache, faster) now. Recommend torchvision first, escalate only if eval shows a
speed/accuracy gap.
254 3. Pose now or parked ? recommend parked behind the interface (box-first).
255 4. Golden labels ? feature-level fusion (P3) is blocked on golden rally labels;
the source is
256 the VLC `rally-annotator` + the Roora vendor pilot
([GOLDEN_SET_VENDORING.md](archives/roora-vendor/GOLDEN_SET_VENDORING.md)).
257
258 ---
259
260 ## 10. Phased execution (owner-gated; ONE PR per slice, off `master`)
261
262 - P0 ? this design doc** + [ROORA_LABELING_GUIDELINES.md](archives/roora-
vendor/ROORA_LABELING_GUIDELINES.md), indexed in
263 `docs/README.md`, cross-linked from `HOW_RALLY_DETECTION_WORKS.md` and ADR-007. *(this
deliverable;
264 docs-only, no code.)*
265 - P1 ? extract the shared person cue:** ? DONE (2026-06-13). Lifted the
torchvision detector +
266 ByteTrack + velocity out of `yolo_hybrid` into
`backend/pipeline/detectors/person_detector.py`
267 (`PersonDetector`); `yolo_hybrid` now orchestrates only and delegates to `self.person`
(no behavior
268 change ? regression-safe, proven by the unchanged process_video output tests). New
269 `tests/test_person_detector.py`; suite +1 green; person_detector is type-clean (left
mypy quarantine
270 for the moved code). The fusion segmenter (P2) reuses this same producer over shuttle-
seeded windows.
271 - P2 ? the FusionSegmenter (Gate A + Gate B + person-feature persistence):** ? DONE
(2026-06-14).**
272 Shipped `backend/pipeline/segmenters/fusion_hybrid.py::FusionSegmenter` (registry
`fusion_hybrid`,
273 default-OFF ? `default_segmenter` unchanged), subclassing the trajectory engine
via two identity
274 hooks (`_enrich_candidate_stats`, `_select_for_handoff` ? wasb/tracknet stay byte-
identical, adversarially
275 verified). It seeds windows from the shuttle substrate, always persists

```

```

`net_crossings`** (Gate A signal),
276     and ? when `indexing.fusion.cue_flags.person` is on ? runs the person cue over each
window's frames
277     (`PersonDetector.detections_per_frame`) and persists the ~5 person features
(`fusion_features.py`,
278     pure + unit-tested). Served gating is optional + default-OFF: `fusion.min_crossings`
(Gate A) /
279     `fusion.min_players` (Gate B) skip sub-threshold windows from the handoff while
**persisting the full
280     superset** (offline substrate intact). Court mask = optional `fusion.court_polygon`
(None ? count-only).
281     Lift to be measured via the harness; held-shuttle-drop + person-feature values are
GPU/owner-side.
282     - **P3 ? feature-level (learned) fusion:** add the person features (incl.
`shuttle_player_colocation`)
283     to `distill_local` / `classifier`; recalibrate IoU on the golden set; report **LOVO
lift on held-out
284     videos**. *(Gated on golden labels.)*
285     - **P4 ? additional cues:** wire **audio** (ADR-007's #1, already scaffolded in
286     `backend/pipeline/audio/hit_detector.py`) and later **pose**, through the same
`RallyCue` interface.
287     Each is add-a-producer, not a refactor.
288
289     ---
290
291     ## 11. Verification
292
293     - **Offline, runnable in this dev container (no GPU/torch/cv2):** unit tests for the
fusion gate logic
294     and `RallyCue` aggregation (court-mask filter, player-count constraint, `? recent-
track` tolerance)
295     and the shuttle-state signals (net-crossing count ? extends the existing `rally_gate`
tests ? and
296     person-colocation) with **synthetic trajectories + person boxes**. A key-license
assertion test
297     (no `ultralytics` import; detector factories stay torchvision/Apache).
298     - **Machine-side (owner, GPU/WSL ? the full pipeline can't run in this container):** run
the
299     `FusionSegmenter` on a real static-tripod clip; confirm **held-shuttle / warm-up /
spectator false
300     positives drop** vs shuttle-only (this should reproduce the owner's observed bug and
show it fixed ?
301     start by confirming Gate-A-in-production alone fixes it); **LOVO IoU lift on the
golden set** for the
302     learned path.
303     - **Docs (P0):** links resolve; both docs indexed in `docs/README.md`; this doc cross-
linked from
304     `HOW_RALLY_DETECTION_WORKS.md` (the fusion open-question) and ADR-007.
305

```

13. user (2026-06-21T13:56:34.259Z)

```

1     # How rally detection works (the 2-layer mental model)
2
3     A short conceptual map of how we go from raw video to rally `[start, end]` segments ?
4     and, crucially, **which layer our annotation can improve.**
5
6     ...
7     video ??? frames ??? [DETECTION] ??? shuttle (x,y,visible) per frame
8         ? ?
9         (WASB CNN, per-frame trajectory
10         frozen weights) ?
11         ?
12         velocity signal (how fast the shuttle moves)
13         ?

```

```

14                                [SEGMENTATION] ??? the 3 knobs we tune
15                                ?
16                                ?
17                                rally windows [start, end] ??? IoU vs golden
labels
18    ...
19
20    ## Two layers ? keep them separate in your head
21
22    **1. DETECTION ? "where is the shuttle, this frame?"**
23    WASB is a CNN that, for each frame, outputs a heatmap ? the shuttle's `(x, y)` + a
24    visibility/confidence. Its weights were learned (by gradient descent, on shuttle-
coordinate
25    labels) and we use them **frozen**. We chose a *shuttle* detector on purpose: "shuttle
in
26    flight" is the signal for "a rally is happening," and it's **camera-agnostic** ? unlike
27    player-motion, which fooled the old heuristic on broadcast cameras that pan with the
players.
28    (Your videos are static tripod, so player-motion would also work ? but shuttle-motion is
the
29    robust general bet.)
30
31    **2. SEGMENTATION ? "when is a rally?"**
32    From the per-frame `(x, y)` we compute the shuttle's **velocity** (displacement between
frames).
33    Then three knobs turn that into rally intervals:
34    - **`velocity_thresh`** ? a frame is "active" only if the shuttle is moving faster than
this
35    (in flight, not lying on the floor / being picked up).
36    - **`merge_gap`** ? active frames within this many seconds get joined into one rally
(bridges
37    brief gaps: a shuttle momentarily occluded, or the beat between two hits).
38    - **`min_window_duration`** ? throw away windows shorter than this (blips / false
flickers).
39
40    Output: a list of `[start, end]` windows = the rallies.
41
42    ## Where the annotation fits (the key point)
43    Our Roora/golden labels are **rally `[start, end]` times** ? they say *when* rallies
are, not
44    *where the shuttle is*. So they can't teach WASB's detector (that needs `(x,y)` labels).
They
45    **can** grade and tune the **segmentation layer**: we score predicted windows against
the golden
46    rallies with **temporal IoU** (how much they overlap), and pick the 3 knobs that
maximise it ?
47    **calibrate on some rallies, measure on held-out ones** so the number is honest.
48
49    ## Where sparks could land (open questions ? not commitments)
50    Things worth a shower-thought; any of these could beat 3 hand-tuned thresholds:
51    - **Richer features than raw velocity:** acceleration, trajectory *continuity* (a real
rally is a
52    smooth arc that survives occlusion), shuttle *height/arc*, or **direction reversals**
(each
53    reversal ? a shot ? ties to `shots_count`!).
54    - **A tiny learned segmenter** instead of 3 thresholds: feed the trajectory features to
a small
55    model trained on the golden rally labels (still the segmentation layer, just smarter).
56    - **Hysteresis** (two thresholds: enter-rally high, stay-in-rally low) ? robust to mid-
rally dips.
57    - **Rhythm/periodicity:** rallies have a hit cadence; dead time doesn't.
58    - **Fusion:** combine shuttle-motion (WASB) with player-motion for static cams. ?
designed in
59    [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (one rally-fusion layer;
the person cue is

```

```

60     ADR-007's parked #2 and `rally_gate.py`'s "Gate B" ? now built (default-OFF,
unmeasured on golden) in `fusion_hybrid.py`).
61     - **Per-video calibration as a *feature*:** "watch a few labeled rallies ? auto-tune to
this
62     court/lighting ? segment the rest sharply" ? exactly your earlier idea, at the right
layer.
63
64     The whole point of per-video calibration: a handful of *your* labeled rallies nudges
these knobs
65     to *your* court/camera, and we report the IoU lift on the rallies it never saw.
66

```

14. assistant (2026-06-21T13:56:37.907Z)

Now let me read the remaining docs: EXPERIMENT_HARNESS, GOLDEN_REGRESSION_FIXTURES, ROORA_LABELING_GUIDELINES, and check QUALITY_ITERATIONS for the latest measured state to validate the \$7 tracker freshness.

15. user (2026-06-21T13:56:38.877Z)

```

1     # Experiment harness ? A/B + shadow eval for rally-detection quality (no-regression by
design)
2
3     **Status:** **P1 IMPLEMENTED** (`backend/eval/experiment.py` +
`tests/test_experiment.py`, shipped
4     2026-06-13) ? the `Variant` layer, `compare()` (labeled A/B, LOVO-honest),
`compare_unlabeled()`
5     (SxS on NEW unlabeled footage), and the leaderboard. **P2?P4 remain owner-gated.**
Phases below are
6     sequenced but each is a separate PR.
7
8     > **Companion docs:** the cues this experiments *on* are in
9     > [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md); the scoring mechanics are
10    > [EVALUATION.md](EVALUATION.md); the no-regression gate + golden labels +
`baseline.json` come from
11    > [GOLDEN_SET_VENDORING.md](archives/roora-vendor/GOLDEN_SET_VENDORING.md); the modular-
cue architecture is **ADR-007**
12    > ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the distilled
classifier is
13    > **ADR-004**. The 2-layer mental model is
[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md).
14
15    ---
16
17    ## 1. Context ? why this, why now
18
19    PR #112 (merged) added the multi-signal fusion design ? shuttle + person + audio cues,
all **default
20    OFF**. The owner's directive for *how* we evolve quality from here:
21
22    > "Evolve the codebase so we can keep finding new ways to improve quality and
experiment with them ?
23    > or add them as a signal ? without suffering any regressions if the idea doesn't work
out. Rolling
24    > back the code would be too dangerous."
25
26    So **experimentation and deployment must be decoupled.** The contract:
27
28    1. Merge experimental code freely ? it's **gated OFF**, so it cannot change production
behavior.
29    2. Test it **offline against golden labels** (most experiments never touch the served
path at all).
30    3. Promote a variant to default **only on measured lift**, through a CI eval gate.
31    4. **"Rollback" = flip a config flag, never `git revert`**. The proven path is a pinned
variant that is

```

```

32         always present.
33
34     **The pleasant surprise (Explore audit, 2026-06-13):** ~80% of this harness already
exists and is
35     simply not composed *as a system*. The missing piece is a thin **variant/experiment
layer + the
36     discipline**, not a platform. §4 is the exact code surface so the next session does not
re-discover it.
37
38     ---
39
40     ## 2. The thesis ? separate expensive DETECTION from cheap DECISION
41
42     The expensive, GPU/WSL-bound, risky part is **detection**: "what signals exist per
frame" (shuttle
43     WASB trajectory, person tracks/boxes, audio hits). The cheap, swappable, pure-Python
part is the
44     **decision**: "given those signals, where are the rallies" (windowing + gates + the
classifier).
45
46     > **Run detection once per video, persist all per-cue signals + per-candidate features,
then re-score
47     > any number of variants OFFLINE ? deterministically, with no GPU and zero production
risk.**
48
49     This is not aspirational ? `distill_local.py` and `calibrate_local.py` already do
exactly this on
50     stored WASB trajectories. The candidate-segments table (`stats` JSON) is the persistence
substrate.
51     ? Most experiments are pure re-scoring of stored data; they never enter the served
pipeline.
52
53     ---
54
55     ## 3. What already exists (the audit ? do NOT rebuild)
56
57     Exact public surface, with `file:line`. **This table is the anti-duplicate-work
artifact** ? start here.
58
59     ### 3.1 Scoring (variant-agnostic ? grades any `List[Interval]`)
60     `backend/eval/metrics.py`
61     - `interval_iou(a, b) -> float` (:15) ? temporal IoU of two `(start,end)` intervals.
62     - `match_intervals(preds, gts, iou_threshold=0.5) -> MatchResult` (:48) ? greedy 1-to-1
IoU matching.
63     - `score(preds, gts, iou_threshold=0.5, match_result=None) -> Score` (:104) ? P/R/F1,
mean IoU, boundary errors; `Score.as_dict()`.
64     - `pr_curve(preds, gts, thresholds=None) -> List[Score]` (:138).
65
66     ### 3.2 A/B + cross-validation primitives (the comparison engine already exists)
67     `backend/eval/calibration.py`
68     - `improvement_report(predict_fn, baseline_config, tuned_config, gts, iou_threshold=0.5)
-> dict` (:148) ? **before/after metrics + %?. This IS the A/B delta.**
69     - `leave_one_video_out(videos, configs, metric="f1", iou_threshold=0.5) -> dict` (:113)
? **LOVO: pick on other videos, score on held-out video (honest cross-video).**
70     - `cross_validate(predict_fn, configs, gts, k=5, metric="recall", iou_threshold=0.5) ->
dict` (:72) ? within-video k-fold overfit guard.
71     - `select_best(predict_fn, configs, gts, metric="f1", iou_threshold=0.5) -> (Config,
float)` (:61).
72     - `evaluate(preds, gts, iou_threshold=0.5) -> Score` (:41); `kfold_indices(n, k)` (:46).
73     - LOVO video dict shape: `{"name": str, "predict_fn": PredictFn, "gts": [Interval]}`.
74
75     ### 3.3 No-regression gate + baselines (the promotion gate already exists)
76     `backend/eval/regression.py`
77     - `regress(db, video_id, ground_truth, stage="final", iou_threshold=0.5, detector=None,
tolerance=DEFAULT_TOLERANCE, baseline_path=DEFAULT_BASELINE_PATH) -> RegressionResult` (:104) ?

```

```

score stored preds vs GT, flag if P **or** R drops > tolerance.
78     - `regress_with_provider(...)` (:160) ? same, fetching GT via the annotation provider.
79     - `save_baseline(video_id, stage, precision, recall, f1, ..., gt_fingerprint=None)`
(:80) . `load_baselines(path) -> dict` (:68).
80     - `RegressionResult` (:43): `regressed`, `reasons`, `gt_hash`, baseline P/R,
`tolerance`, `has_baseline`; `as_dict()`.
81     - Default baseline file: `eval_baselines/regression_baseline.json` (:33).
82
83     `run_regression.py` ? CI-gatable CLI, `main(argv)` (:52). **Exit codes: 0 ok / 1
regression / 2 no-DB / 3 no-baseline.** Flags: `--video-id --tier --annotations
--stage{candidates,final} --detector --iou --tolerance --baseline --update-baseline --allow-
missing-baseline --json`.
84
85     ### 3.4 Pinnable model artifact
86     `backend/eval/classifier.py` ? `LogisticRegression` (pure numpy, L2):
`fit(X,y,feature_names=,class_weight="balanced")` (:40), `predict_proba` (:75),
`predict(threshold=0.5)` (:81), `to_dict`/`from_dict` (:86/:97), `save`/`load` (JSON)
(:107/:111). **A trained model = one JSON file ? a pinnable, hashable variant artifact.**
87
88     ### 3.5 Offline re-scoring substrate (persist once, re-score forever)
89     `backend/infrastructure/database.py`
90     - `add_candidate_segment(video_id, detector, start_time, end_time, chunk_index=None,
confidence=None, stats=None, detection_params=None) -> Optional[int]` (:318) ?
`stats`/`detection_params` stored as JSON; `candidate_segments` table cols incl. `stats`,
`detection_params`, `human_label`, `confirmed_segment_id` (schema :97).
91     - `query_candidate_segments(video_id, detector=None, only_unlabeled=False) ->
List[dict]` (:355) ? returns rows with deserialized `stats`.
92
93     ### 3.6 Label assignment + feature glue
94     `backend/eval/training.py`
95     - `label_candidates(candidate_intervals, gt_intervals, iou_threshold=0.5) -> List[int]`
(:50) ? y-labels by IoU match to golden (same matcher as the evaluator).
96     - `build_training_set(candidate_rows, gt_intervals, iou_threshold=0.5,
feature_fn=extract_yolo_features) -> (X, y, intervals)` (:64) ? pluggable `feature_fn`.
97
98     ### 3.7 Existing LOVO drivers to mirror (not duplicate)
99     - `backend/eval/distill_local.py` ? `run(specs, iou=0.5, threshold=0.5, model_out=None)`
(:98), CLI `python -m backend.eval.distill_local --manifest videos.json --model-out ...`;
`FEATURE_NAMES = [duration, peak_velocity, mean_velocity, active_ratio, net_crossings]` (:40);
already does LOVO over Gemini silver labels.
100    - `backend/eval/calibrate_local.py` ? `run(specs, iou=0.5, metric="f1")` (:77);
`LocalConfig = (velocity_thresh, merge_gap, min_window_duration, min_crossings)`; grid + LOVO.
101    - `backend/eval/calibrate_wash.py` ? `run(...)` , `load_golden_gts()`;
`backend/eval/gt_loader.py::load_gt_intervals(csv_path, *, strict=True)` is THE canonical golden
reader.
102
103    ### 3.8 Registries + config knobs (variant selection)
104    - `backend/pipeline/segmenters/base.py` ? `SegmenterRegistry` (:35), singleton
`segmenter_registry` (:63); registered (6): `motion`, `gemini`, `yolo_hybrid`,
`tracknet_hybrid`, `wasb_hybrid`, `fusion_hybrid`.
105    - `backend/annotations/base.py` ? `AnnotationProviderRegistry`, singleton
`annotation_registry`; providers `file`, `auto`.
106    - `backend/config/models.py::IndexingConfig` ? `default_segmenter` (:65, default
`"yolo_hybrid"`), `detector_model` (:72), `detector_score_threshold` (:73),
`yolo_velocity_threshold` (:74), `yolo_frame_skip` (:75), `min_segment_duration` (:68),
`dead_time_merge_gap` (:69), `motion_threshold` (:67). The `min_crossings` knob now lives on
`FusionConfig` (default `0` = OFF; also `min_players`, both served Gate A/B thresholds ? see
fusion P2).
107
108    ### 3.9 Confirmed ABSENT (no duplication risk)
109    Explore found **no** existing experiment / variant / shadow / A-B *machinery* ? only a
stray "A/B test"
110    aspiration comment in `backend/pipeline/detectors/base.py` and an "experiment E1" note
in
111    `archives/research/AUDIO_RALLY_DETECTION_PLAN.md`. The `regress()` + baseline path is

```

```

the canonical comparison mechanism.
112
113     ---
114
115     ## 4. The variant abstraction (the one thin thing to add)
116
117     A Variant = a declarative recipe, not a code branch:
118
119     ```
120     Variant = (segmenter_name, IndexingConfig overrides, enabled cues + per-cue params,
121               classifier model path/hash)
122     ```
123     JSON-serializable; selected via the existing `segmenter_registry` +
124     `IndexingConfig.default_segmenter`.
125     Control = a *pinned, frozen* variant (recipe + model hash) that is always registered
126     and is the
127     default. Adding a new cue = a new Variant differing by one flag ? the control recipe is
128     untouched.
129
130     ---
131
132     ## 5. Three experiment modes
133
134     ### 5.1 Offline A/B (primary, label-driven, zero production risk)
135     `compare(videos, variant_A, variant_B)` ? a thin driver composing existing
136     functions:
137     `query_candidate_segments()` (load persisted features) ? re-score each variant ?
138     `calibration.improvement_report()`
139     + `calibration.leave_one_video_out()` + `metrics.score()`, graded vs golden
140     (`gt_loader.load_gt_intervals`).
141     Reports per-video and LOVO-honest aggregate IoU/P/R/F1 deltas. This is the A/B
142     test, and it needs
143     almost no new scoring code ? only the glue.
144
145     Honest reporting (default ON, additive): compare(..., honest_stats=True) also
146     attaches a "honest"
147     block ? per-metric bootstrap CIs, a paired ?F1 CI + exact sign test (C1), and
148     F1@k +
149     segment-count ratio (C4) ? alongside the bare-mean fields (existing output
150     unchanged; honest_stats=False
151     restores the legacy shape). A bare LOVO mean at n?6 is not promotable on its own; the
152     lift is real when the
153     ?F1 CI clears 0 and the sign test agrees. Wiring of the course corrections in
154     [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) §13/C1+C4 (helpers:
155     `backend/eval/eval_stats.py`,
156     `segmentation_metrics.py`); record_experiment carries the honest ?F1 into the
157     leaderboard.
158
159     ### 5.2 Shadow (for cues that touch detection, e.g. the person detector)
160     The new cue runs and persists its signal into the candidate
161     `stats`/`detection_params`, but the
162     served decision stays control. Offline you then ask "what would variant B have
163     scored?" via §5.1 ?
164     B never affects output until it wins. This is how a detection-touching change (not just
165     a re-scoring
166     change) earns its way in without risk.
167
168     ### 5.3 Promotion gate (no-regression, CI-enforced)
169     regression.regress() + baseline.json + run_regression.py exit codes. A variant
170     becomes the new
171     default only if it doesn't regress the golden set beyond tolerance (exit 1
172     blocks). Control is
173     pinned ? rollback = flip default_segmenter/variant config, never a code revert.
174
175
176

```



```

157     ---
158
159     ## 6. No-regression guarantees (the heart of the owner's ask)
160
161     - **New cues/signals default OFF** ? per-cue enable flag in `IndexingConfig`; merged
code can't change
162         behavior until explicitly enabled.
163     - **Control variant pinned** ? frozen recipe + classifier model hash, always registered
& default.
164     - **Promotion only via the eval gate** ? `run_regression.py` in CI (exit 1 blocks); no
silent default change.
165     - **Experiments are records** ? variant-spec hash ? per-video + LOVO scores + label-set
hash + timestamp,
166         persisted to a small JSON **experiment leaderboard** (generalizes `baseline.json`).
Answers "which
167         signals ever helped, on which video slices," and makes every result reproducible.
168     - **Decoupled events** ? "merge an experiment" and "change the default" are separate,
independently
169         revertible actions.
170
171     ---
172
173     ## 7. What to build later (gap list ? thin, additive, owner-gated)
174
175     1. ~**`backend/eval/experiment.py`** ? a `Variant` dataclass + `compare(videos, a, b)`
composing the $3
176         functions (no new scoring math).~ **DONE (2026-06-13).** Shipped `Variant` (+ pinned
`CONTROL`),
177         `compare` (labeled A/B with LOVO-honest learned-variant training, mirroring
`distill_local`),
178         `compare_unlabeled` (the **SxS-on-new-videos** mode ? agreed / A-only / B-only via
`match_intervals`,
179         no GT), `record_experiment` (? `eval_baselines/experiment_leaderboard.json`), and
`load_video_candidates`
180         (the `query_candidate_segments` bridge). JSON variant recipes first as recommended; a
`VariantRegistry`
181         only if the count grows. 19 tests incl. the no-regression invariant.
182     2. **Persist per-cue features** into candidate `stats` (person: `players_in_court`,
`two_sided_motion`,
183         `in_court_ratio`, `shuttle_player_colocation`; shuttle-state: `net_crossings`,
`direction_reversals`;
184         audio: hit times) so any future variant re-scores offline ? ties to
`MULTI_SIGNAL_FUSION_PLAN.md` $3.2/$4.
185     3. **Per-cue enable flags** in `IndexingConfig` (default `False`); `default_segmenter`
already exists.
186         Add the `min_crossings` knob (fusion P2) here too.
187     4. **CI promotion lane** using `run_regression.py` against the golden `baseline.json`;
write the
188         experiment-leaderboard JSON.
189
190     ---
191
192     ## 8. Phasing (owner-gated; ONE PR per slice, off `master`)
193
194     - **P0 ? this doc**, indexed in `docs/README.md`, cross-linked from
`MULTI_SIGNAL_FUSION_PLAN.md` and
195         `GOLDEN_SET_VENDORING.md`. *(this deliverable; docs-only, no code.)*
196     - **P1 ? `experiment.py` `compare()`** ? **DONE (2026-06-13):** thin wrapper over
`calibration` +
197         `metrics` + `query_candidate_segments`; pure-logic, fully tested in this dev container
(`compare`,
198         `compare_unlabeled`, `record_experiment`, the no-regression invariant).
`compare_unlabeled` adds the
199         SxS read on **new unlabeled footage** the owner asked for (run control + candidate on
a fresh video,

```

```

200     diff where they agree/diverge ? two reels via the existing `VideoCompiler` is the GPU-
side follow-on).
201     - **P2 ? per-cue enable flags + feature persistence** into candidate `stats` (default
OFF) ? couples
202         with fusion P1/P2.
203     - **P3 ? CI promotion gate**: `run_regression.py` lane on the golden `baseline.json`;
experiment-leaderboard JSON.
204
205     ---
206
207     ## 9. Scope guardrail
208
209     **Single-tenant / local-first** (per `PLATFORM_ARCHITECTURE.md`): a JSON experiment
registry + the
210     existing harness ? explicitly **NOT** MLflow / W&B / a service. Keep it the size of
`baseline.json`.
211
212     ---
213
214     ## 10. Verification
215
216     - **Offline, runnable in this dev container (no GPU/torch needed for the glue):** unit
test that
217         `compare()` is deterministic on synthetic candidate rows; a **no-regression invariant
test** asserting
218         "all new-cue flags OFF ? predictions byte-identical to control"; a CI-gate test
asserting
219         `run_regression.py` exit codes (0 ok / 1 regress / 3 no-baseline).
220     - **Machine-side (owner, GPU/WSL):** persist candidates once on a real clip, run
`compare(control,
221     fusion)` LOVO, confirm the reported lift matches a manual `metrics.score()` spot-
check.
222     - **Docs (P0):** links resolve; indexed in `docs/README.md`; cross-linked from
223         `MULTI_SIGNAL_FUSION_PLAN.md` and `GOLDEN_SET_VENDORING.md`.
224     </content>
225

```

16. user (2026-06-21T13:56:39.818Z)

```

1     # Golden Regression Fixtures ? video-free, non-attributable regression suite
2
3     > **Status:** `IN PROGRESS` ? owner-gated . **Owner:** Avi . **Created:** 2026-06-16 .
**Last updated:** 2026-06-16 . **Legal:** owner **accepted the researched risk** 2026-06-16
(condition §3 stand; this is the owner's informed acceptance, **not** a retained-counsel
opinion)
4     > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE, per archive policy)
5     > **Tracking anchors:** §7 progress tracker is the source of truth; indexed in
`docs/README.md`; pointer in memory `current-state` (and `docs/NEXT_STEPS.md` once work starts).
6     > **Relation to existing docs:** EXTENDS (does not supersede)
`docs/GOLDEN_DATA_SHARING.md` (#175); reuses the eval stack documented in `docs/EVALUATION.md` /
`docs/EXPERIMENT_HARNESS.md`; quality floor ties to `docs/QUALITY_ITERATIONS.md`.
7     > **Honesty note:** §3 (Legal) is **researched general information, NOT legal advice** ?
it gates on retained counsel (§8). §2/§4/§7 mark which claims are **[verified]** against code vs
**[design]** proposals.
8
9     ---
10
11     ## 0. TL;DR
12
13     Let a contributor on **another machine with no video, GPU, WSL, or DB** run the rally-
segmentation regression suite from **committed structured files**, so refactors and non-quality
PRs get a red/green side-effect signal. Two layers, built once and grown by accretion:
14
15     1. **Mechanics** ? freeze the deterministic pipeline at the **post-detector trajectory

```

```

seam** and replay it through the *already-shipped* eval stack (`experiment.py` / `metrics.py` /
`windowing.py` / `regression.py`), with two auto-discovered gates: a
**characterization/snapshot** gate (catches any behaviour change) and a **quality-floor** gate
(catches quality drops).
16     2. **Privacy/legal filter (SEALED-FIXTURES)** ? only **de-attributed, owner-source,
biometric-free** numbers are ever committed; richer/sensitive streams stay key-gated or out-of-
repo. Non-attributability comes from **deletion-at-source**, not encryption. ? "Minimized" means
**stream minimization** (shuttle-only, no person/audio/pose) ? **NOT** reduction of the shuttle
stream to ~5 scalars: the de-attributed per-frame `trajectory.csv` IS committed (see §4c + its
risk-acceptance re-confirmation note).
17
18     **The convergence that makes this clean:** the *privacy* answer (publish shuttle-only,
drop person/pose streams) and the *determinism* answer (freeze only the `CONTROL` path ? which
runs **no numpy execution**, keeping the numpy/BLAS learned scorer off the cross-machine gate)
point at the **same minimal public tier**. That coincidence is the spine of the design. (numpy
is still *imported* transitively via `experiment` and must be installed; the CONTROL replay just
never calls into it.)
19
20     **Two findings already actionable, independent of the rest:**
21     - ? **Live de-attribution bug** *(verified)*: real player/venue names are committed
today in `eval_baselines/heuristic_lovo_n6.json`,
`eval_baselines/gemini_wrapper_2026-06-10.json`, and `backend/eval/eval_partial_labels.py` (the
last also has a personal `C:/Users/avidu/...` path) ? and in git history. ? **Deliverable P0.**
22     - ? **Determinism trap** *(verified, corroborated by two independent red-teams)*: the
learned variants (`shuttle_only`, `fusion`) train a numpy/OpenBLAS `DYNAMIC_ARCH` logistic
regression whose decisions flip near the 0.5 threshold across CPUs ? would false-RED on a
contributor's machine. ? the cross-machine gate must be **CONTROL/static-only**.
23
24     ---
25
26     ## 1. Problem & goal
27
28     The rally-quality measurement is split across two machines by capability
(`docs/GOLDEN_DATA_SHARING.md`):
29     - `backend/eval/fusion_golden.py` **extracts** `_golden_features.json` ? needs the
**video + trajectory CSV + labels + GPU/WSL** (GPU box only).
30     - `backend/eval/fusion_compare.py` / `backend/eval/ablation.py` **re-score** those JSONs
on **pure CPU, no video** ? any machine.
31
32     The CPU re-scoring is *already* video-free, but the inputs live in gitignored `output/`
/ private OneDrive, so the regression tests that depend on them **skip everywhere**
(`tests/test_ablation.py::test_litmus_reproduces_gate_a`, `tests/test_fusion_compare.py`). The
committed `eval_baselines/` baselines exist but `regression.py` still scores DB-stored
predictions (needs a pipeline run on the video).
33
34     **Goal:** commit a tiny, version-tagged, **numeric-only, non-attributable** per-fixture
corpus + auto-discovered gates so a clean `git clone` / CI gets a real red/green regression
signal with **zero video/GPU/DB**, while never leaking attributable or sensitive data.
35
36     ---
37
38     ## 2. Decisions locked
39
40     | # | Decision | Source | Implication |
41     |---|-----|-----|-----|
42     | D1 | **Corpus is owner-recorded / owned** | owner answer 2026-06-16 | Legal analysis
takes **Fork A** (clean): no third-party copyright/ToS infringement; concern is trade-
secret/moat + privacy. Strongest non-attributability (no public reference to correlate against).
|
43     | D2 | **Repo private now, may open-source later** | owner answer 2026-06-16 | Design
for **public-grade hygiene from day one** ? git history is permanent; nothing scrubable-later. |
44     | D3 | **Adults, consent not formalized** | owner answer 2026-06-16 | **Person/pose/gait
streams stay OUT of the public tier** (DPDP/GDPR derived-personal-data risk without consent).
Public tier = **shuttle + abstract intervals + non-biometric features only**. |
45     | D4 | **Freeze seam = post-detector trajectory** *(design; SHIPPED differs ? see ?)* |

```

workflow-1 synthesis | Widest deterministic surface; reuse the shipped eval stack. ?

****Correction (audit 2026-06-16):**** the shipped freeze ****commits** the de-attributed per-frame `(Frame,Visibility,X,Y)` `trajectory.csv` ****** (shuttle-only, time-origin reset) ? it is ****NOT**** minimized to ~5 aggregate scalars. See §4c (and its risk-acceptance re-confirmation note). |

46 | D5 | ****Cross-machine gate is CONTROL/static-only**** *(design)* | both red-teams [verified cause] | Learned variants (numpy/BLAS) are ****owner-box/Tier-1 only****, asserted with tolerance, never as a cross-machine byte-snapshot. |

47 | D6 | ****Non-attributability = deletion-at-source, encryption = public-only barrier**** *(design)* | workflow-2 synthesis | A contributor who runs the tests necessarily reads plaintext+key (DRM/analog-hole). So we ****delete**** attribution, not hide it. Encryption only stops the no-key public. |

48 | D7 | ****Two tiers**** *(design)* | workflow-2 synthesis | ****Tier-0**** public/no-key (shuttle-only, de-attributed). ****Tier-1**** key-gated/out-of-repo (person/audio, full corpus), ****skips cleanly when absent****. |

49

50 ---

51

52 ## 3. Legal foundation (researched ? gates on counsel, §8)

53

54 ****Core verdict (settled across IN/US/EU):**** the derived numbers ? per-frame shuttle (x,y), rally start/end intervals, audio/court/aggregate-person feature values ? ****do NOT inherit** the source video's copyright. ****** They are uncopyrightable *facts/measurements*, not authored expression, and not a §101 "derivative work."

55

56 | Jurisdiction | Copyright of the data | Database right | Contract / ToS | Net |

57 |---|---|---|---|---|

58 | ****India**** (primary) | Not copyrightable as facts (*EBC v. D.B. Modak*); but compilations get thin protection on a low "skill/labour" bar (*Burlington*, *OLX v. Padawan*) ? only bites if ****lifted from a 3rd-party dataset****, not self-derived | ****None**** (no sui-generis right) ? a real permissiveness edge | Anti-scraping ToS (untested in court) + ****IT Act s.43/s.66**** civil/criminal exposure for 3rd-party ingestion | ****Owned-source + minors-excluded + no person/pose = clean**** |

59 | ****US**** | Not copyrightable (*Feist*; §102(b); *NBA v. Motorola*) | None (*Feist* killed sweat-of-the-brow) | ****Dominant risk.**** ToS enforceable where copyright/CFAA fail (*ProCD*, *hiQ* ? injunction + ****forced deletion****). Remedy for redistribution is takedown, not nominal damages | Publishable as facts ****iff**** owned/permissioned source + no barring ToS + no person/biometric/minor stream |

60 | ****EU / UK**** | Not copyrightable (*Football Dataco v. Yahoo*) | ****Genuinely contestable.**** *Football Dataco v. Sportradar*: recording pre-existing on-court facts = ****obtaining****, not "creating" ? our "we created it" premise is ****overconfident****; harm test = *CV-Online Latvia* | Ryanair v. PR Aviation: ToS can ban extraction/redistribution even with no IP right | Get EU+UK advice before redistributing numeric streams at scale |

61

62 ****Provenance is the decisive fork.**** ****Fork A (owner-recorded, unpublished)**** ? the only path safe to commit publicly; collapses to a trade-secret *choice* (Indian breach-of-confidence / Contract Act 1872 / IT Act s.72 ? ****not**** US DTSA) + privacy. ****Fork B (broadcast/YouTube/scraped)**** ? do ****not**** commit; the extraction act needs a license or now-****contested**** US fair use (*Thomson Reuters v. Ross* 2025; USCO May-2025), and unlawful acquisition independently sinks it (*Bartz v. Anthropic*, ~\$1.5B settlement). Caveat: owner footage ****shot at a third party's event**** (tournament/club) carries ticket/venue/organiser/broadcaster terms ? treat as Fork B until cleared.

63

64 ****Publish-by-stream (privacy):****

65 - ? ****Publish-OK:**** shuttle (x,y)+Visibility (inanimate object), abstract rally/GT ****intervals**** (identity-stripped), non-biometric audio/court/aggregate-person features detached from identity.

66 - ?? ****Conditional (effectively never in v1):**** person bounding boxes that *persistently single out a player* ? flips to GDPR Art.9 the moment it identifies. This architecture is prone to that flip ? treat per-player tracking as a tripwire.

67 - ? ****Never publish:**** pose/skeleton/****gait**** (re-identifies at 80?95% even with faces blurred ? *Weiss et al. 2025*), face crops, any named-player?movement mapping. Matches the existing "gait/biometrics strictly future (GDPR Art.9)" guardrail.

68 - ? ****Minors flip everything:**** DPDP "child" = under 18; s.9(3) ****prohibits behavioural monitoring**** (the Rule-12 exemptions don't cover this project). A persistent per-player signature on a minor is, conservatively, prohibited. ****Exclude minors**** from any persisted

stream beyond inanimate shuttle/interval data.

69

70 ****Non-attributability is relative, not absolute.**** A per-video trajectory is a fingerprint (de Montjoye: 4 points ? 95% re-ID). It's only non-attributable because the ****Fork-A source stays unpublished**** (no public reference to correlate against). Per ***EDPS v. SRB*** (CJEU C-413/23 P, 2025) it ****remains personal data to the holder**** who can re-run the detector. So: conditional, revocable (publish the source and the link snaps back), forward-looking (re-assess as the public-video environment grows). Never market it as "anonymous."

71

72 ***Key authorities:** Feist 499 U.S. 340; NBA v. Motorola 105 F.3d 841; 17 U.S.C. §101/§102(b); Sega v. Accolade 977 F.2d 1510; Football Dataco v. Yahoo (C-604/10) & v. Sportradar (C-173/11); BHB v. William Hill (C-203/02); CV-Online Latvia (C-762/19); Ryanair v. PR Aviation (C-30/14); Thomson Reuters v. Ross (D. Del. 2025); EBC v. D.B. Modak; DPDP Act 2023 s.3/s.9 + Rules 2025; EDPS v. SRB (C-413/23 P). Full URLs in the research artifact (§10).

73

74 ---

75

76 **## 4. Architecture (design)**

77

78 **### 4a. Freeze seam + reuse map**

79 Freeze at the ****trajectory CSV ? CONTROL-path replay****. One shared helper ``backend/eval/golden_fixtures.py::replay_fixture()`` is the ***single*** code path used by both the freeze tool and both gates. As ****shipped****, that path is (per the module docstring):

80 ``parse_trajectory_csv ? distill_local.build_candidates(proposal=(vt,mg,mwd)) ? VideoCandidates(5 shuttle cols only) ? experiment.predict(CONTROL) [keep all ? merge_overlaps] ? metrics.score``.

81 Reuses ****verbatim****: ``metrics.score/interval_iou`, `experiment.predict/CONTROL/VideoCandidates`, `distill_local.build_candidates`, `tracknet_runner.parse_trajectory_csv`/`trajectory_to_action_windows`, `rally_gate` net-crossings (inside `build_candidates`), `gt_loader.load_gt_intervals` (strict), `regression.gt_hash`, `eval_baselines/` committed-baseline precedent.`

82 > ? ****Correction (audit 2026-06-16):**** the replay does ****NOT**** import ``windowing.py`` ? there is no ``windowing.build_windows(preset)`` / ``WindowingPreset`` / "default SERVED" call. The windowing params travel as a ****raw 3-tuple `(vt, mg, mwd)`**** pinned in each fixture's ``provenance.json`` and passed straight to ``build_candidates``; the harness deliberately avoids the named ``windowing.SERVED`` constant so a future SERVED retune can't silently churn fixtures. (``segmentation_metrics``, ``experiment.compare/Variant.spec_hash``, ``regression.save_baseline``, ``eval_stats`` are ****not**** on the M1 replay path either ? the quality-floor gate M4 will add the baseline/stats pieces.)

83

84 **### 4b. Tiered model (SEALED-FIXTURES)**

85 | Tier | Where | Key? | Contents | Test behaviour |

86 |---|---|---|---|---|

87 | ****Tier-0**** | ``eval_baselines/fixtures/`` (tracked) | none | opaque surrogate slug, fps/frame_width (retained, NOT quantized), the ****de-attributed per-frame `trajectory.csv`**** ``(Frame,Visibility,X,Y)`` (shuttle-only, time-origin reset), the 5 shuttle feature columns inside ``expected.json``, strict-slug ``label``, relative ``gts``. ****No person/audio/pose.**** ? NOT minimized to 5 scalars ? see §4c. | ****Unconditional**** ? runs on every PR/fork; turns today's skipped litmus into a real public gate |

88 | ****Tier-1**** | OneDrive (default) ****or**** SOPS-encrypted in-repo | token/age key | full corpus ****with**** person/audio; authoritative ?F1/CI/sign-test + exact Gate-A litmus | ****Skips cleanly**** when key/data absent (generalize ``_golden_present()`` ? ``_tier1_present()``) |

89

90 A fresh no-key clone is always ****green****; coverage never depends on a secret.

91

92 **### 4c. Anti-attribution measures (Tier-0)**

93

94 > ? ****RISK-ACCEPTANCE RE-CONFIRMATION REQUIRED (audit 2026-06-16).**** An earlier draft of this doc claimed Tier-0 commits "only ~5 aggregate shuttle scalars" and **"**never**** the per-frame CSV". ****That is not what ships.**** The shipped freeze (``freeze_real_fixture``) ****commits the FULL per-frame `(Frame,Visibility,X,Y)` `trajectory.csv`**** ? de-attributed (time-origin reset + opaque random slug + metadata strip), but ****not minimized to 5 scalars****. The per-frame stream is the de Montjoye fingerprint (4 points ? 95% re-ID against a published reference). The owner's 2026-06-16 risk acceptance (§7 P2) was recorded against the inaccurate "5 scalars" basis and ****MUST be RE-CONFIRMED on this corrected basis**** before any ****real**** Tier-0 fixture is committed

(P3). Synthetic fixtures (current state) carry zero provenance risk regardless.

95

96 - ****Opaque RANDOM surrogate id**** (`os.urandom` + a private surrogate?source map kept off-git) ? ***not*** the MD5 `video\_id`, and ***not*** HMAC(salt, video_id?name) (red-team: brute-forceable over a tiny known roster if the salt leaks). Never persist the source `name`/filename.

97 - ****Strict-slug `label` + dropped `source\_note`**** ? the operator `label` is validated against `^[a-z0-9][a-z0-9\_\-]\*\$` (a free-text identifier like a player/event name is refused), and operator free-text `source\_note` is ****accepted-but-NOT-persisted**** (never reaches the committed `provenance.json`) ? both close the free-text identifier leak `\_scan\_forbidden` could not police (audit fix #1).

98 - ****Metadata strip**** ? no filename, video_id, match/event/player identity, capture date.

99 - ****Time-origin reset**** ? subtract per-fixture `t0` (6dp-rounded) from all start/end/gts. ****Metric-invariant within `\_FLOAT\_TOL`**** (`metrics.interval\_iou` and start/end-error are pure differences): the genuine (raw) reset-vs-no-reset score delta sits at the float-error floor (~3.3e-7); bit-identical for whole-second offsets. See §4c-note below and §6.

100 - ****Per-frame trajectory IS committed (de-attributed), NOT minimized to scalars**** ? Tier-0 commits the de-attributed per-frame `trajectory.csv` (shuttle stream only). It is ****not**** reduced to ~5 aggregate scalars; the 5 shuttle `feature` columns appear (computed) inside `expected.json`, but the raw per-frame `(Frame,Visibility,X,Y)` stream is part of the committed fixture. (This is the corrected basis the ? note above gates on.)

101 - ****Stream minimization (streams, not resolution)**** ? Tier-0 keeps the shuttle block only; person/audio excluded; pose/gait/face have ****no field in the schema by design**** (positive `\_scan\_forbidden` assertion on the written bytes).

102 - ****fps/frame_width are RETAINED as scale constants**** ? they are NOT quantized today (quantization is a deferred P3 option, see §7 ?). Reconcile any "quantize fps/frame_width" wording elsewhere to this.

103 - ****Provenance hard-fail gate**** ? the de-attribution pass refuses to emit a public fixture for any record not tagged owner-recorded(Fork-A) + minors-excluded + non-blank-license, and refuses a negative reset label (timeline predates the trajectory). Writes go to a temp dir and are promoted only after the privacy scan passes, so a refused freeze leaves no partial fixture (audit fixes #3a/#3b).

104

105 **### 4d. The two gates**

106 - ****Characterization/snapshot**** ? replay fixture, ****parse-compare**** JSON (never byte-compare ? CRLF/float-format safe; `core.autocrlf=true` here), exact on int/structural fields, `approx(abs=1e-6)` on floats; stamp + check windowing-preset/label/gt fingerprints (incommensurable ? loud fail). Catches any behaviour change even at constant F1. ****CONTROL path only**** across machines.

107 - ****Quality-floor**** ? score vs committed labels; assert per-video + mean F1 ? floor in `eval\_baselines/regression\_baseline.json` (created here; carries a ****windowing fingerprint**** so a served-default change is flagged incommensurable, not silently compared); emit `eval\_stats` paired-delta-CI + sign test (NUMBERS INFORM); add a ****metrics-vs-base-branch**** paired delta. Synthetic-tier floors measure ****correctness, not field quality**** (documented).

108

109 **### 4e. Serialization & optional protobuf**

110 JSON with `sort\_keys=True` + pre-quantized floats + LF is sufficient. ****Protobuf is optional (01)**** ? adopt only if cross-OS byte-stable snapshots are wanted (proto3 `rally.fixture.v1`, codegen + `protoc && git diff --exit-code` guard, `\*.pb` binary -text`). Protobuf is **\*\*encoding, not secrecy\*\*** (trivially `--decode_raw`).

111

112 **### 4f. Encryption & keys (Tier-1 only)**

113 ****Default = out-of-repo + token**** (existing OneDrive `RALLY\_GOLDEN\_DIR` seam ? zero sensitive bytes in git history, sidesteps the legal "redistribution" act). ****Alternative = SOPS + age**** (per-recipient, revocable, value-level diffs) for Fork-A-clean-but-private data wanting in-repo versioning. Rejected: git-crypt (whole-repo re-key), git-lfs (storage ? secrecy). CI secret exposed ****only** to the Tier-1 job on trusted branches, never to fork-PR runners. The ****HMAC/surrogate salt**** (if used) and the surrogate?source map live with the private corpus, rotated per release. Document explicitly that Tier-1 crypto defends ****only**** the no-key public.

114

115 ---

116

117 **## 5. Threat model (who's protected vs not)**

118

119 | Actor | Protected against | NOT protected (by design / residual) |

120 |---|---|---|

```

121 | **TM1 no-key public** | map fixture?source/match/player; decrypt Tier-1; confirm a
source even holding the candidate file | read Tier-0 numbers + run shuttle-only harness
(intended; useless-in-isolation) |
122 | **TM2 keyed contributor** | recover filename/video_id/match/face/pose/gait ? *never
written* (survives the analog hole) | read every value the suite reads (unpreventable; not
pretended otherwise) |
123 | **TM3 CI runner** | leaked log can't leak attribution it never had | runs the suite
(intended); Tier-1 secret never on fork runners |
124 | **TM4 malicious keyed insider** | correlation needs a public reference; Fork-A source
is unpublished ? "not reasonably likely" | **residual:** conditional + revocable ? publish/leak
the source later and the link can snap back (forward-looking re-test duty) |
125 | **TM5 re-ID adversary** | source unpublished + reset timeline ? no back-match against
a public reference | ? becomes live if a Fork-A source is ever published ? **and the committed
per-frame trajectory IS a fingerprint** (de Montjoye 4-point re-ID), NOT just ~5 scalars
(corrected §4c). The protection is the unpublished source, not minimization |
126 | **TM6 accidental plaintext commit** | required CI leak-scan rejects MD5/known-
name/non-surrogate/plaintext-SOPS | if the required check is disabled, a leak lands |
127
128 ---
129
130 ## 6. Honest limits ? what a green suite does NOT prove
131
132 - Green proves only that the **deterministic post-trajectory transforms** are unchanged
for the **frozen cases**. It does **not** cover: the WASB/TrackNet detector,
`_probe_fps_width/cv2` frame-peek, the AI/provider handoff, chunking/re-encode, the ffmpeg
stitch, DB persistence/migrations, or the API filter. *A refactor that breaks decoding or the
detector passes these gates green.* The owner-box `regression.py` DB run remains the only full-
chain check.
133 - Characterization = **behaviour-locking, not correctness** ? it freezes current
behaviour *including current bugs*. "Passing" = "behaviour unchanged", never "behaviour
correct". This caveat must ride in the **test pass/fail message**, not only the docs.
134 - **Synthetic-tier floors = plumbing correctness, NOT field quality.** Never cite
synthetic F1 in `QUALITY_ITERATIONS.md` / the paper as rally-detection performance (tag `kind:
synthetic`).
135 - Tier-0 public coverage is **shuttle-only** ; the person/audio fusion ?F1 and exact
Gate-A values need Tier-1 (skipped on no-key clones).
136 - Data is **not "anonymous"*** absolutely (EDPS v. SRB) ? relative, conditional,
revocable, forward-looking.
137 - The committed Tier-0 stream is the **de-attributed per-frame `trajectory.csv`**, NOT
~5 scalars ? a per-frame trajectory is a de Montjoye fingerprint; the only protection is the
**unpublished source** (§4c ?, §5 TM5).
138 - **Time-origin reset is metric-invariant within `_FLOAT_TOL`, not byte-exact for every
offset.** The genuine (raw) reset-vs-no-reset score delta is at the float-error floor (~3.3e-7);
it is **bit-identical for whole-second offsets**, but because 6dp-rounding is not additive
(`round6(a)-round6(b) ? round6(a-b)`) the *serialized* 6dp boundary error can differ by one 6dp
ULP (1e-6) on a sub-second offset ? the serialization grid, not a metric divergence. (Audit fix
#2 rounds `t0`+labels to 6dp before the shift to keep the residual ? the grid step.)
139 - Encryption gives **no secrecy from a running contributor** ; protobuf gives **no
secrecy at all**.
140 - De-attribution does **not** cure provenance/licensing ? it sits *behind* the per-
record provenance gate + counsel sign-off.
141 - `--refresh-snapshot` can rubber-stamp a real regression; mitigation is **review of the
visible diff** + a separate, deliberate floor-baseline update. **A recorded `--reason` is
REQUIRED** (the CLI errors without one ? audit fix #6).
142
143 ---
144
145 ## 7. Deliverables & progress tracker ? **source of truth**
146
147 Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. One **small PR per row**,
branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.
148
149 | ID | Deliverable | Depends on | Counsel-gated? | Status | PR |
150 |---|-----|-----|-----|-----|---|
151 | **P0** | **Leak remediation + required CI leak-scan.** Scrub real names/paths in

```

```

`eval_baselines/heuristic_lovo_n6.json`, `gemini_wrapper_2026-06-10.json`,
`backend/eval/eval_partial_labels.py`, `tests/test_cleanup_caches.py` ? opaque surrogate keys +
private map; add CI leak-scan (MD5 / known-name / non-surrogate / plaintext-SOPS) as a
**required status check** ; decide on git history (PR #77, accept vs `filter-repo`). | ? | No
(owned data hygiene) | ? | ? |
152 | **M1** | **Fixture framework + synthetic Tier-0 + shared replay helper** (also
delivered M2-core + M3 ? see below). `backend/eval/golden_fixtures.py` (`replay_fixture`, schema
dispatch, `Path(__file__)`-relative) + 3 **tool-generated** synthetic fixtures + manifest +
README + `.gitattributes` LF pin + `tests/test_golden_fixtures.py`. CONTROL-only (no numpy
execution on the replay path); cross-OS verified. | ? | No (synthetic) | ? |
[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) |
153 | **M2** | **Freeze/refresh extractor.** Core (`freeze_synthetic` / `refresh_snapshot` /
`--check` CLI + idempotence test) shipped in M1 (#189); the `--license`/`--minors-
free`/`--owned` **refusal gate** shipped with P1 (this PR). | M1 | No | ? |
[#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) . _this PR_ |
154 | **M3** | **Characterization gate** ? `tests/test_golden_fixtures.py`: parse-compared
(exact-int / approx-float 1e-6), **CONTROL-only**, positive no-person/audio assertion,
perturbation + idempotence. **Shipped in M1 (#189).** (Windowing/label fingerprint guard + row-
count cap ? fold into M4.) | M1 | No | ? | [#189](https://github.com/avidullu/badminton-
highlight-indexer/pull/189) |
155 | **MX** | **Cross-OS determinism guard (CI).** Lightweight `golden-crossos` CI job
(ubuntu/windows/macros) running `--check` + the fixtures test ? locks the cross-OS guarantee
(CONTROL replay is numpy-only, no GPU stack; a *separate fast job*, not full-suite×3). | M1 | No
| ? | [#191](https://github.com/avidullu/badminton-highlight-indexer/pull/191) |
156 | **M4** | **Quality-floor gate + first `regression_baseline.json`.**
`tests/test_golden_quality_floor.py`; create `eval_baselines/regression_baseline.json` with
**windowing fingerprint** ; metrics-vs-base paired-delta; corpus-level (not per-slug) LOVO for
runtime. | M1, M2 | No | ? | ? |
157 | **P1** | **De-attribution + minimization** ? `golden_fixtures.freeze_real_fixture()`
(extends the M1 module, not a separate script): **opaque random `fx_hex` surrogate slug**,
metadata strip, metric-invariant **time-origin reset**, strict GT loader, **positive no-
person/audio/MD5/source-path assertion**, manifest non-clobber; `--freeze-real` CLI with the
**refusal gate** (exit 2 unless owned + minors-free + license). Tooling-only, 12 synthetic
tests; **no real data committed**. | M1 | No (tooling) | ? | _this PR_ |
158 | **P2** | **Owner-accepted (researched-risk), 2026-06-16.** Owner accepted the
researched IP/privacy analysis (§3) to proceed **under its conditions** (owned/Fork-A source .
minors excluded . biometric/person streams excluded . de-attributed). Owner's informed risk
acceptance ? **NOT a retained-counsel opinion** ; optional future counsel confirmation per §8. |
P1 | Owner-accepted | ? | ? |
159 | **P3** | **Tier-0 REAL subset.** Commit de-attributed, minors/biometric-excluded,
owned Fork-A real fixtures (unblocked by P2). Hard-gated **in code** by P1's refusal gate (owned
+ minors-free + license required) ? needs the owner's per-video clearance list. | P1, P2 |
conditions §3 | ? | ? |
160 | **O1** | *(optional)* **Protobuf + deterministic serialization.** `rally.fixture.v1`
proto3, codegen + `protoc`/`git diff --exit-code` guard, canonical writer, `.gitattributes`
binary. Only if cross-OS byte-stable snapshots are wanted. | M1 | No | ? | ? |
161 | **O2** | *(optional)* **Tier-1 key-gated path.** `scripts/fetch_golden.py` (OneDrive
token, default) and/or SOPS+age; `_tier1_present()` ; `RALLY_GOLDEN_DIR` wired; Tier-1 CI job
gated on secret, never on fork runners. | M1 | No | ? | ? |
162 | **DOC** | **Discoverability cross-links** ? added the `eval_baselines/fixtures/` row
to `CODE_MAP §0` + a `golden_fixtures.py` entry to `CODE_MAP §2.3` ; cross-linked
`GOLDEN_DATA_SHARING.md` (committed subset ? private OneDrive corpus). README indexed in #186;
memory `current-state` kept live. (EVALUATION/QUALITY_ITERATIONS cross-links **deferred** ? that
doc surface is hot in the concurrent quality track; avoid collision.) | M1 | No | ? | _this PR_
|
163
164 | **Must-fix-before-any-public-commit (red-team), folded into the rows above:** ? numeric-
tolerance not byte-equality + multi-OS matrix + CONTROL-only (M3/D5/MX ?) ; ? scrub names in
**all** files + history decision (P0) ; ? leak-scan a **required** check, hooks are convenience
only (P0) ; ? person-optional loader + **positive** no-person assertion, don't trust the silent
`0.0` default (P1 ?) ; ? random surrogate id not HMAC-over-roster (P1 ?) ; ? quantize/drop
fps/frame_width (P1 ? fps/fw retained as scale constants; quantize is a P3 option) ; ? document
Tier-0 omits the fusion path (DOC) ; ? counsel + Fork-A + minors/biometric exclusion as hard
preconditions (P2 owner-accepted; enforced in code by P1's refusal gate ?).
165

```



```

166     ### Progress log (quantitative ? coverage must hold/improve, CI incl. the 3-OS cross-OS
job green)
167     | PR | Deliverable(s) | Suite | Coverage | Notes |
168     |----|-----|-----|-----|-----|
169     | [#189](https://github.com/avidullu/badminton-highlight-indexer/pull/189) | M1 +
M2-core + M3 | 746 pass / 7 skip | 80.41% | synthetic Tier-0 + CONTROL replay +
characterization; cross-OS proven (CI Linux replayed Windows-frozen fixtures) |
170     | [#191](https://github.com/avidullu/badminton-highlight-indexer/pull/191) | MX +
tracker/legal reconcile | full suite green | 80.41% | `golden-crossos` CI job green on
ubuntu+windows+macos ? cross-OS **enforced**; legal ? owner-accepted |
171     | _this PR_ | P1 + M2 refusal gate | 759 pass / 7 skip | **80.49%** (+0.08) |
`freeze_real_fixture` + refusal gate; **+13 tests**; tooling-only, no real data committed |
172     | _this PR (audit-hardening)_ | adversarial-audit fixes #1?#17 | 784 pass / 7 skip |
**80.71%** (+0.22) | strict-slug `label` + dropped `source_note` (#1); 6dp-rounded reset within
`_FLOAT_TOL` (#2); temp-dir-then-promote + negative-label refusal (#3); exact-membership source-
path scan (#4); key-set-equality in `compare_expected` (#5); `--reason` mandatory (#6);
independent schema guard (#7); `max_win` recorded-not-applied note (#8); numpy-execution wording
(#9); +failure-semantic & by-kind tests (#10/#11); `test_golden_real_fixtures.py` added to
`golden-crossos` CI (#12); doc reconciled to shipped reality incl. per-frame-CSV honesty + risk
re-confirm (#13?#17). Synthetic `--check` byte-identical (unchanged). |
173
174     ---
175
176     ## 8. Open owner / counsel questions
177
178     **Counsel questions ? documented basis (the owner accepted the researched risk on
2026-06-16; see P2). These remain the record, and a retained-counsel confirmation is
optional/recommended before scaling beyond a small owned, minors-free, de-attributed subset:**
179     1. Does the EU/UK **sui-generis database right** subsist in our self-derived
shuttle/rally streams on the *Sportradar* "obtained vs created" reasoning, and does
redistributing a de-attributed Tier-0 subset cause *CV-Online* harm?
180     2. Post *Thomson Reuters v. Ross* + USCO-2025, is **US fair use for the extraction
step** defensible if the dataset could compete in a "data/analytics" market ? does publishing
only numbers (never footage) matter?
181     3. India: are anti-scraping ToS enforceable; **IT Act s.43/s.66** exposure for any 3rd-
party ingestion; does pending *ANI v. OpenAI* (~Apr 2026) move it? Confirm **DPDP s.9/Rule 10**
and whether per-player tracking is prohibited s.9(3) monitoring of a minor.
182     4. Sign-off on the **final committed Tier-0 subset + provenance log**; is non-
attributability-via-private-source a defensible posture to assert given its
conditional/revocable nature?
183
184     **Owner (strategic):**
185     5. Publishing the Fork-A subset **forfeits the trade-secret moat** (Indian breach-of-
confidence framing) ? weighed against community/benchmark/paper aims. Your call.
186     6. **Forward-looking commitment:** do not publish a highlight reel of the *same* source
whose fixtures are committed (without re-clearing), and accept a periodic re-assessment as the
public-video environment grows.
187     7. **Scope of first public commit:** synthetic-only (zero provenance risk, recommended)
vs a counsel-cleared real subset.
188     8. **Tier-1 mechanism:** out-of-repo+token (recommended) vs SOPS+age in-repo. And: is
protobuf (01) wanted, or is canonical-JSON enough?
189
190     ---
191
192     ## 9. Definition of done
193
194     The project is **DONE** (flip `Status ? DONE`, archive) when:
195     - [ ] **P0, M1?M4, P1, DOC** are ? and merged.
196     - [ ] A clean `git clone` on **win + ubuntu + mac** runs `python run_all_tests.py` and
the **Tier-0 gates pass green with no key / no video / no GPU / no DB**.
197     - [ ] The **CI leak-scan is a required status check** and passes (no MD5 / real-name /
non-surrogate / plaintext-SOPS in the tree).
198     - [ ] Characterization + quality-floor gates are **parse-compared, CONTROL-only,
windowing-fingerprinted**, and demonstrably catch a seeded behaviour change.
199     - [ ] `docs/CODE_MAP.md` §0 + TEST MAP updated; this doc cross-linked; `NEXT_STEPS.md` +

```

```

memory `current-state` point here.
200     - [ ] Either the public commit is synthetic-only, or (P2 owner-accepted ? +
P3) for a real-data commit, under the §3 conditions enforced in code by P1's refusal gate.
201     - [ ] §6 honest-limits caveat rides in the test pass/fail messages, not only the docs.
202
203     *Optional 01/02 may remain ? without blocking DONE (they're earned, not assumed).*
204
205     ---
206
207     ## 10. References
208
209     Internal code anchors (verified): `backend/eval/{regression.py,fusion_golden.py,fusi
on_compare.py,ablation.py,experiment.py,metrics.py>windowing.py,distill_local.py,gt_loader.py,ev
al_stats.py}` · `backend/pipeline/detectors/{base.py,tracknet_runner.py,rally_gate.py}` ·
`backend/utils/hashing.py::compute_video_id` · `eval_baselines/{heuristic_lovo_n6.json,gemini_wr
apper_2026-06-10.json,experiment_leaderboard.json}` ·
`tests/test_ablation.py::test_litmus_reproduces_gate_a` · `.github/workflows/ci.yml`.
210     Internal docs: `GOLDEN_DATA_SHARING.md` (#175), `EVALUATION.md`,
`EXPERIMENT_HARNESS.md`, `QUALITY_ITERATIONS.md`, `MULTI_SIGNAL_FUSION_PLAN.md`, `CODE_MAP.md
§0`.
211     Legal authorities: see §3 (full URLs in the research workflow artifact). This
section is researched general information, not legal advice.
212
213     ---
214
215     ### Changelog
216     - 2026-06-16 ? Initial draft synthesized from two multi-agent workflows (mechanics
design + 2 critiques; legal memo IN/US/EU + SEALED-FIXTURES architecture + red-team). Verified
the live name-leak and the BLAS-determinism cause against code. Status: DRAFT, owner-gated.
217

```

17. assistant (2026-06-21T13:56:44.132Z)

Now let me read the ROORA labeling guidelines (it's in archives per the cross-links) and check QUALITY_ITERATIONS for the latest measured state to validate tracker freshness.

18. user (2026-06-21T13:56:46.433Z)

```

GOLDEN_SET_VENDORING.md
README.md
ROORA_LABELING_GUIDELINES.md
---QUALITY ITER head---

```

Rally-detection quality iterations ? living log

Purpose. A running, reverse-chronological log of how we evolve rally-detection quality: each iteration's hypothesis, what shipped, the measured (or expected) effect, and the lesson. This is the

live continuation of the historical trail in

[archives/quality-iterations/](archives/quality-iterations/README.md) ? per the archive policy (`CLAUDE.md`), **terminal-state** (DONE/REJECTED) work gets a dated subdir there; this file stays at `docs/` top level and keeps growing, then its entries graduate to the archive when they conclude.

Methodology (non-negotiable). Every quality change rides the experiment harness (`EXPERIMENT_HARNESS.md`)(`EXPERIMENT_HARNESS.md`): new cues land **flag-gated**, default OFF*; they're graded **offline against golden labels** (per-video + LOVO-honest); promotion to default happens **only** on measured lift* through the `run_regression.py` gate; **rollback** = flip a config flag, never `git revert`. Cues are designed in `MULTI_SIGNAL_FUSION_PLAN.md`(`MULTI_SIGNAL_FUSION_PLAN.md`); scoring in `EVALUATION.md`(`EVALUATION.md`); the 2-layer model in `HOW_RALLY_DETECTION_WORKS.md`(`HOW_RALLY_DETECTION_WORKS.md`).

****Golden corpus = the 15 labelled matches**** in `output/human\_lovo\_manifest.json` (the original 6 ?
`adarsh\_avi\_singles, kushagra\_singles, targetest\_doubles, gbaaddy, testlarge\_short, mahadevpura\_singles` ? owner-confirmed 2026-06-13, since grown to 15 with the later GoPro/Boxhill clips). The model trained on these is what we run alongside the existing detector on NEW footage for the side-by-side review.

Baselines that drive strategy (IoU 0.5 vs human golden labels)

Path	Mahadevpura (clean-ish)	Kushagra (multi-court)
Heuristic (velocity windowing)	0.657	0.157
Gen-0 owned head (LOVO, n=6)	0.744	0.561
Two-stage WASB+Gemini refine	0.83	?
Gemini wrapper (`gemini-flash-latest`)	0.93	0.66

****Reading:**** clean footage is commoditized (serve via Gemini); ****multi-court** drops the commodity wrapper to 0.66 ? that is the owned model's ship target. ****Baselines** tracked in `eval\_baselines/`. The owner's observed false positive ? a player ****holding/flicking**** the shuttle between points counted as a rally ? is the precision bug the fusion gates target.

Current track (live) ? newest first

P7 ? Long-rally (>5s) quality lens (eval) + reel-selection knob ? MEASURED, default-OFF *(2026-06-18)*

****Hypothesis (owner).**** Local OVER-FIRES, and its false positives are mostly SHORT (motion blips, background-court fragments on multicourt). Filtering the eval to ****long rallies (>5s)**** ? the eye-catching highlights that carry the product ? should strip most short FPs and reveal local's real long-rally quality, answering the North-star "do we beat Gemini on the single-match long-rally happy path?". ****Duration**** (end?start), NOT shot-count, is the filter: golden `shots\_count` is blank and the no-AI path reports it "Unknown", whereas duration comes from the start/end labels we already trust.

****What shipped (both default-OFF, zero behaviour change at the defaults).****

- ****Eval lens**** ? `rally\_seg\_eval.py --min-duration <s>` filters BOTH predictions AND golden GTs to rallies longer than `s` *before* matching, consistently across the tIoU and R2 tolerance paths (a FILTER, not a new metric). Plus `--stratify` / `--long-tau` (default 5s): a `{all vs long} × {single vs multicourt}` aggregate table. Court type = the committed `MULTICOURT\_CLIPS` set (the manifest is gitignored, so the set is the auditable source), overridable by a manifest `court\_type` field; kept in sync with [`GOLDEN\_VIDEOS.md`](data_pipeline/GOLDEN_VIDEOS.md).
- ****Product knob**** ? `stitching.min\_rally\_duration` (s, 0 = OFF) drops short rally pairs before the reel compile in `tools/generate\_highlights.py` ? a cleaner reel-selection filter than `min\_shot\_count` (which leans on the "Unknown" shot_count).

****Measured (n=15 golden, milestone cap6+R4, LOCAL vs golden; reproduce: `python -m backend.eval.rally\_seg\_eval --stratify`):****

stratum	court	n	tolF1	tolP	tolR	tIoU F1@0.5	seg_ratio
---------	-------	---	-------	------	------	-------------	-----------

```
|---|---|---|---|---|---|---|
| all | single | 11 | 0.375 | 0.326 | 0.464 | 0.482 | 1.434 |
| all | multicourt | 4 | 0.294 | 0.238 | 0.393 | 0.430 | 1.601 |
| all | all | 15 | 0.353 | 0.302 | 0.445 | 0.468 | 1.478 |
| **long>5s** | single | 11 | 0.264 | 0.278 | 0.267 | 0.416 | **1.037** |
| **long>5s** | multicourt | 4 | 0.240 | 0.197 | 0.310 | 0.430 | **1.597** |
| **long>5s** | all | 15 | 0.257 | 0.256 | 0.279 | 0.420 | 1.186 |
```

Pooled counts: single **608 preds / 416 GTs** (only 36% of preds are >5s vs 50% of GTs ? long pred:GT ? **218:210 = 1.04**); multicourt **288 preds / 179 GTs** (55% of preds >5s ? long pred:GT = **158:100 = 1.58**).

Findings ? the lens turned a one-line hypothesis into a 3-part measured map (NUMBERS INFORM):
 1. ? **Single court, count-level: thesis HOLDS.** Local's over-production is dominated by SHORT windows
 ? restricting to >5s collapses seg_ratio **1.43 ? 1.04** (preds ? GTs in count). The short-blip FPs
 are real and duration-filterable.

19. assistant (2026-06-21T13:56:57.690Z)

This confirms the corpus has grown to 15 (the \$7 tracker in RALLY_QUALITY_RESEARCH still says n=6 in places). Let me read the R00RA guidelines briefly and check the rest of QUALITY_ITERATIONS to capture the latest measured state and any newer items beyond the tracker.

20. user (2026-06-21T13:56:58.636Z)

```
1      # Roora Annotation Guidelines ? Rally Golden Set
2
3      **Version: v8 (2026-06-13).** Supersedes the older Drive copies (v3/v5/v6/v7) **and**
the internal
4      "DRAFT v2" label ? those numbers never matched each other; **v8 is the single label to
use with
5      Roora going forward.** Content is otherwise the finalized 2026-06-12 guideline plus the
6      fusion-driven additions marked **`? ADDED (fusion)`** below.
7
8      > **?? Canonical home / provenance.** **This in-repo file is the sole canonical copy of
the guideline**
9      > ? it lives here because the multi-signal-fusion features that motivate the new fields
(court
10     > calibration, singles/doubles) originate in *this* repo (see
11     > [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) and
12     > [GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md), ADR-006). The annotation
deliverable still imports
13     > into the sibling **`sports-data-collector`** repo via `python -m src.cli.curate
import-local`; if a
14     > mirror is ever staged there for the collector's intake docs, it must be regenerated
from this file
15     > (and the `? ADDED (fusion)` callouts are exactly the deltas to port).
16     >
17     > **`TODO(avidullu)` markers are preserved** ? they are owner-only items (Drive links,
intake
18     > manifest, NDA counsel review) and must be filled before sending to Roora. Search for
`TODO(avidullu)`.
19
20     Brief (vendor spec) + Illustrated Guide (rater how-to), in one doc.
21
22     ---
23
24     # PART A ? Roora Annotation Brief (vendor-facing spec)
25
26     **Project:** racket-sport rally segmentation golden / regression dataset.
```

```

27     **Sports in scope:** Badminton, Pickleball, Table Tennis (each in singles AND doubles).
28
29     This brief is the vendor-facing spec. The annotation deliverable is a per-rally CSV that
imports
30     directly via `python -m src.cli.curate import-local`.
31
32     ## 0. What we need in one sentence
33
34     For each provided match video, produce ONE CSV file listing EVERY rally with its start
time, end
35     time, number of shots, and how it ended ? EXCLUDING all dead time.
36
37     ## 1. Deliverable format (canonical)
38
39     One CSV per video (UTF-8), named exactly `.csv` (we supply the `video_id`).
One row per
40     rally, chronological, header row required.
41
42     **Columns:**
43
44     | Column | Type / values | Meaning |
45     |---|---|---|
46     | `rally_number` | integer, 1-based, chronological ? no gaps, no repeats | rally index |
47     | `start_mmss` | `m:ss.mmm` (e.g. `1:12.40`) | rater-entered serve-contact time |
48     | `end_mmss` | `m:ss.mmm` | rater-entered ball/shuttle-dead time |
49     | `start_time` | decimal **seconds** (e.g. `72.40`) ? **AUTO-FILLED** from `start_mmss`
| machine time |
50     | `end_time` | decimal **seconds** ? AUTO-FILLED from `end_mmss` | machine time |
51     | `shots_count` | integer | total racket/paddle/bat contacts (serve = shot 1) |
52     | `ending_reason` | `winner` \| `forced_error` \| `unforced_error` \| `service_fault` \|
`let` \| `other` | WHY it ended |
53     | `last_shot_outcome` | `in` \| `net` \| `out` \| `n_a` | WHERE the last shot landed
(`n_a` for service_fault/let) |
54     | `score_before` | optional, "A-B" ? blank if not confidently readable | score before
|
55     | `score_after` | optional, "A-B" | score after |
56     | `notes` | optional free text ? **REQUIRED when `ending_reason = other`** | edge-case
note |
57
58     **Times (mm:ss ? decimal seconds):** raters type only `start_mmss` / `end_mmss`. The
template Google
59     Sheet fills `start_time` / `end_time` with this formula (shown for `start_time`, row 2):
60

```

21. user (2026-06-21T13:56:59.813Z)

```

80     are real and duration-filterable.
81     2. ?? **Multicourt: thesis FAILS.** Even among >5s rallies local over-fires **1.58x** ?
adjacent-court
82     games are genuine LONG rallies ? long *spurious* windows. **A duration filter does
NOT clean
83     multicourt; that needs court-localization** (Gate-B / court mask ? gap-closure Lever
C). So "FPs are
84     mostly short" is a *single-court* statement.
85     3. ?? **Strict tolerance precision does NOT jump** (tolP single 0.326 ? 0.278) even
where the count
86     normalizes, because the >5s pred set is contaminated by **over-MERGED windows**
(kushagra med 9.9s,
87     mahadevpura_1's 63s blob, GX010141 14.8s) and cap=6 truncates genuinely-long rallies
so the strict
88     |?end| ? 1.5 s match fails (tIoU, forgiving, holds ~0.42; strict tol drops). The
long-rally
89     *strict-eval* happy path is gated by merge-control + the cap-vs-long tension, not
just FP-stripping.
90     4. ? **Product knob is still a clean reel win.** A reel doesn't need strict boundary

```

tolerance ? an

91 over-merged 15 s clip is a fine highlight, and dropping <5 s pairs removes the single-court short

92 junk. Defaults OFF; owner reviews the numbers before any flip.

93

94 ****Gemini:**** not computable in-checkout ? the Boxhill gemini intervals live in the owner-side scorer's DB

95 (not committed). The lens is metric-agnostic: feed Gemini's intervals through the same

96 ``filter_by_min_duration`` to get its long-rally row (owner-side one-liner). From gap-closure §3, Gemini

97 already fires fewer/longer windows with ~0 short FPs, so a long filter mostly affects LOCAL.

98

99 ****Lesson.**** A plausible product one-liner ("filter to long ? precision jumps") split into three measured

100 truths once instrumented: single-court count ?, multicourt ? (needs a court signal), strict-precision ?

101 (needs merge control). Default-OFF lens + knob; the owner decides any flip on the numbers. *(P7 in

102 [``RALLY_DETECTION_GAP_CLOSURE.md``](RALLY_DETECTION_GAP_CLOSURE.md) §7.)*

103

104 **### R1 ? MILESTONE LAUNCHED** (cap=6 + R4 flipped default-ON) ? re-verified at n=13 *(2026-06-18)*

105 The R1 milestone is ****shipped**** (PR #204): the W1 cap=6 + R4 inactivity-end defaults are now ON.

106 Full quality-impact table on both rulers, the over-seg guard, and the honest caveats are recorded in

107 the ****[R1 Milestone Launch Report]**(R1_MILESTONE_LAUNCH_REPORT.md)****** (per the Launch-Report convention).

108

109 ****Re-verified on the grown n=13 corpus**** (the 2 new Boxhill clips GX010137/GX010141) ? the milestone

110 holds and the headline is ***stronger*** than at n=11:

111

112 | transition | n=11 (filed) | ****n=13 (launched)**** |

113 |---|---|---|

114 | baseline ? milestone tolF1 | ?+0.232, 10/0, p=0.002 | ****?+0.222, 12/0/1, p=0.0005**** |

115 | cap12 ? cap6 tolF1 | ?+0.108, 10/0, p=0.002 | ****?+0.110, 12/0/1, p=0.0005**** |

116 | cap6 ? +R4 tolF1 | ?+0.047, 8/0, p=0.008 | ?+0.040, ****9/1/3****, p=0.0215 |

117 | net over-seg guard | PASS (mr 0.694?0.150) | ****PASS**** (mr 0.651?0.161, 0 regressions) |

118

119 Both new clips improve (GX010137 0.448?0.685, GX010141 0.333?0.431); tIoU ladder baseline 0.236 ?

120 milestone 0.474. Independently reproduced by a 2nd harness (#214 ``lovo_resweep.py``, matches to 3 dp;

121 LOVO picks cap=6 ****13/13**** on tIoU) and survived a 5-lens adversarial workflow (every L00 fold clears

122 the bar; worst-fold drop GX010128 still ?+0.202, p=0.0010 ? no single clip drives it). ****Honest notes:****

123 W1 cap is the dominant lever / R4 is the marginal +0.040 increment (1 loss = GX010137, a tolerance

124 artifact); the gate is the NET rate-based guard (strict=False by design); GX010141 is the worst-behaved

125 clip (future over-merge target). The nightly self-flags STALE post-flip ? re-baseline once with

126 ``--update-baseline``. W2 / hard-cap / R5 stay default-OFF.

127

128 **### R5 ? blob-fallback for the continuously-active blob** (default-OFF, stratum tool) *(2026-06-17)*

129 The one clip the cap6+R4 milestone still fails is the ****continuously-active blob**** (mahadevpura_1:

130 ~174 s of play with no inactivity gap to split on). The W1 gap-split finds no gap; R4 finds no slow

131 tail; the result is one ~66 s mega-window that matches nothing (****tolF1 0.000****, the

```

degenerate
132 0-match outcome [ `REAL_FOOTAGE_VALIDATION_REPORT.md` ](REAL_FOOTAGE_VALIDATION_REPORT.md)
$5 R5 names).
133
134 **R5** ( `window_blob_fallback`, default-OFF) is the last resort: **after** the gap-split
and the R4
135 trim have each had their chance, any group still longer than `window_blob_factor` × the
cap is a
136 genuine blob ? midpoint-chop it to ? the cap, **log the forced cut**, and **flag every
resulting
137 window `forced_cut=True`** (surfacing the clip as one that needs rally-STATE cues ? net-
crossings /
138 two-sided motion ? not a bigger cap). `_blob_fallback_chop` in `tracknet_runner`, wired
through
139 WindowingPreset / IndexingConfig / the served path / `rally_seg_eval --blob-fallback`.
140
141 **Why a factor gate, and why not just `window_hard_cap`:** `hard_cap` chops **before**
R4 and fires on
142 *every* gap-less over-cap window ? it over-segments normal clips (mahadevpura_2 40?68).
R5 runs **after**
143 R4, and the magnitude gate (**default 4x**) leaves a normal long rally (~1?2× the cap)
alone, force-cutting
144 **only** the degenerate blob. That makes it surgical and do-no-harm.
145
146 **MEASURED (n=11, on the cap6+R4 milestone, R2 `tolF1`):**
147
148 | clip | milestone | + R5 | windows (milestone ? +R5) |
149 |---|---|---|---|
150 | **mahadevpura_1 (the blob)** | **0.000** | **0.238** | 2 (max 65.8 s) ? 32 (max 4.1
s), all `forced_cut` |
151 | mahadevpura_2 (also continuously-active) | 0.122 | 0.180 | 42 ? 100 |
152 | kushagra_singles | 0.078 | 0.110 | 45 ? 71 |
153 | Badminton_BXH_2 | 0.328 | 0.324 | 78 ? 92 (?0.004, noise) |
154 | *(the other 7 clips)* | ? | *unchanged* | no-op (no group exceeds 4× the cap) |
155
156 At factor=4, R5 touches **only** the 2 continuously-active clips + a noise-level BXH2
?0.004 ? effectively
157 **do-no-harm**. It is a **STRATUM tool**: the win is the blob rescue (0.000?0.238,
removing the
158 mean-swinging 0-match pathology), not a corpus-mean lift (agg `tolF1` 0.323?0.354, sign
**3/3/5, n.s.** at
159 factor 4; the factor sweep is monotonic ? 6× is strictly do-no-harm, 3× over-chops).
**Default-OFF**; turn
160 it on selectively for continuously-active footage. The R1 net over-seg guard confirms
it: milestone?+R5 net
161 cost 0.423?0.289, **PASS**, no merge_rate regression. [numbers:
`scratch/r5_evidence.py`]
162
163 ### R1 ? milestone evidence (cap=6 + R4) + the net over-seg promotion guard
*(2026-06-17)*
164 The R1 milestone is the smallest safe local-windowing flip: **W1 bounded windowing cap=6
+ the R4
165 rally-END inactivity trim, no W2** (the per-clip-best config, settled by R3/R4). This
entry (a) re-runs
166 the full ablation ladder baseline?milestone under the **R2** `tolF1` on **n=11** with
honest paired
167 sign-tests, and (b) ships the **net over-seg promotion guard** that unblocks the
milestone.
168
169 **The full ladder (n=11, R2 `tolF1`, asymmetric ?_start=2/?_end=1.5):**

```

22. assistant (2026-06-21T13:57:05.726Z)

I now have a comprehensive picture. Let me verify the \$7 tracker freshness against current code/state ? specifically whether R2/R4 work is reflected in RALLY_QUALITY_RESEARCH.md \$7 (it

appears the §7 table still shows W2 net-crossings as the "next" item and n=6, while R1/R2/R4/R5/P7 have all shipped per QUALITY_ITERATIONS and the launch reports). Let me do a quick check on the gap-closure doc and confirm there are no other live tracker rows.

23. user (2026-06-21T13:57:06.506Z)

```
1      # Rally Detection ? Closing the Accuracy Gap (golden-eval-driven, per-video)
2
3      > **Status:** `DRAFT` ? implementation **owner-gated** . **Owner:** Avi . **Created:**
2026-06-18 . **Last updated:** 2026-06-18
4      > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
5      > **Tracking anchors:** §7 progress tracker is the source of truth; pointer in memory
`current-state`; add to `docs/README.md` + `docs/NEXT_STEPS.md` when work starts.
6      > **Relation to existing docs:** **peer-of**
[ `RALLY_QUALITY_RESEARCH.md` ] (RALLY_QUALITY_RESEARCH.md) (that is the research + eval-framework
doc; this is the *execution tracker* for the remaining accuracy gap). **Builds on**
[ `MULTI_SIGNAL_FUSION_PLAN.md` ] (MULTI_SIGNAL_FUSION_PLAN.md) (signal mechanics),
[ `QUALITY_ITERATIONS.md` ] (QUALITY_ITERATIONS.md) (the R-log / empirical backbone),
[ `GOLDEN_VIDEOS.md` ] (data_pipeline/GOLDEN_VIDEOS.md) (the corpus). Does **not** duplicate them ?
it points.
7      > **Honesty note:** each claim tagged `[verified]` (checked vs code/measured) /
`[design]` (proposed). The headline numbers below are real measurements, but `n` is small ? read
§6.
8
9      ---
10
11     ## 0. TL;DR
12     We are using the **growing golden corpus** (human rally labels) to find and close the
rally-detection accuracy gap **per video, before launch** ? and the first ground-truth clip
already reshaped the picture. The boundary work (R2/R3/R4) **holds** ? when local fires on a
real rally, its edges are tight ( $|?| < 1$  s). **The remaining gap is DETECTION, not boundaries:**
(G1) **recall** ? both local and Gemini miss ~half the rallies on hard clips; (G2) **precision**
? local over-fires (false positives), especially on continuous drilling; (G3) **rally-END over-
extension** ? when a player *quickly knocks the shuttle back after a point*, shuttle-velocity
alone can't tell "rally" from "casual reset," so the window runs long / merges. n=2 sharpens the
priority: **LOCAL = high recall / low precision, GEMINI = high precision / low recall**
(opposite, complementary ? on the normal clip local's recall *beats* Gemini's). So **the #1
lever is precision** (cut local's false positives via `rally_gate` Gate-A + colocation; the
player-state rally-END rule feeds this), with **signal fusion** (local's recall + Gemini-style
precision) as the durable win. **Caveat:** two golden clips so far (n=15 corpus,
GX010141+GX010137) ? directional, not conclusive.
13
14     ## 1. Problem & goal
15     The detector is good enough to ship a highlights product, but the golden eval shows
clip-dependent accuracy that we want to **close consistently on every video, not just on
average** ? catching the failure modes in the quality loop instead of after launch.
16
17     - **Concrete target:** on every corpus video, local is **at-par-or-better than Gemini vs
the golden labels**, with **no per-video regression** on either eval set (the dual-eval-set
discipline, [[eval-dual-set-regression]]).
18     - **"Good" =** the three gaps below closed enough that a milestone flip (default-ON)
clears the over-seg guard AND improves per-video F1 broadly, recorded in a Launch Report
([[launch-report-convention]]).
19     - **Read first to get current:**
[ `RALLY_QUALITY_RESEARCH.md` ] (RALLY_QUALITY_RESEARCH.md) §1?2 (the gap + diagnosis),
[ `MULTI_SIGNAL_FUSION_PLAN.md` ] (MULTI_SIGNAL_FUSION_PLAN.md) §4 (in-flight vs in-hand),
[ `HOW_RALLY_DETECTION_WORKS.md` ] (HOW_RALLY_DETECTION_WORKS.md) (the 2-layer model).
20
21     ## 2. Decisions locked
22     | # | Decision | Source / date | Implication |
23     |---|-----|-----|-----|
24     | D1 | The gap is **DETECTION (recall + precision + rally-END), not boundary precision**
| golden GX010141, 2026-06-18 `[verified]` | Don't build a generic boundary head; spend on
rally-state signals. |
```


25 | D2 | Measure ****per-video AND aggregate; no regression on ANY video**** | [[eval-dual-set-regression]] / [[decision-rigor-norm]] | A win must broadly help, not trade one clip for another. |

26 | D3 | ****Golden labels = ground truth; Gemini = reference, not truth**** | n=15 corpus | Gemini under-detects too ? agreement ? accuracy. Report both, gate on golden. |

27 | D4 | New signals land ****default-OFF + flag-gated****, promoted only on measured lift | EXPERIMENT_HARNESS.md / [[weak-signals-retained]] | Zero-regression rollout; retained even on a small-n null. |

28 | D5 | ****Rally-END needs a PLAYER-state signal**** ? shuttle velocity can't separate a competitive rally from a casual post-point return | owner insight + R4 limitation, 2026-06-18 [design] | Add person kinematics / shuttle-in-hand as the rally-end discriminator. |

29

30 ## 3. Foundation ? the measured gaps (golden ground truth)

31 Ground-truth-scored Boxhill clips (more land as the owner labels them). Milestone windowing = cap6 + R4 ([`QUALITY\_ITERATIONS.md`](QUALITY_ITERATIONS.md) R3/R4). Local & Gemini vs truth, n=2 so far: [verified]

32

33 | clip (golden) | system | found | tIoU@0.5 P/R/F1 | tolF1 | hits | false+ | missed | \?start\ | \?end\ |

34 |---|---|---|---|---|---|---|---|---|

35 | ****GX010141**** (29, hard/drilling) | LOCAL | 22 | 0.59/0.45/****0.51**** | 0.43 | 13 | 11 | 16 | 0.78 | 0.68 |

36 | | GEMINI | 16 | 0.94/0.52/****0.67**** | 0.53 | 15 | 4 | 14 | 0.31 | 0.46 |

37 | ****GX010137**** (34, normal match) | LOCAL | 39 | 0.77/****0.88****/****0.82**** | 0.68 | 30 | 14 | 9 | 0.39 | 0.44 |

38 | | GEMINI | 27 | 1.00/0.79/****0.89**** | 0.89 | 27 | 0 | 7 | 0.50 | 0.52 |

39

40 - ****Corpus aggregate (n=13, milestone):**** tIoU F1@0.5 ****0.474**** [0.359, 0.579], F1@0.25 0.633; R2 tolF1 0.359. GX010137 (0.822) is the corpus's ****best**** clip; GX010141 (0.510) is above the mean. [verified]

41 - ****Opposite failure modes ? the central finding (n=2).**** ****LOCAL = high recall, low precision; GEMINI = high precision, low recall.**** On the normal clip GX010137, local's ****recall 0.88 BEATS Gemini's 0.79**** (local finds **more** real rallies) but local fires 14 false positives; Gemini fires ****zero**** false positives but misses 7 rallies. They are ****complementary**** (on GX010141 they overlap on only ~3 rallies yet each hits ~13?15 of the golden set). ? ****fusion (local's recall + Gemini-style precision filtering) is the clear win.****

42 - ****The gaps, re-prioritized by the data:****

43 - ****G2 ? precision is LOCAL's #1 gap**** (over-fires on both clips: 11/22 and 14/39 false positives). This is the cheapest + highest-leverage fix (Gate-A / colocation). Filter local's FPs on GX010137 and it lands ~0.88+ ? at-par with Gemini, since local already out-recalls it.

44 - ****G1 ? recall**** is clip-dependent: a local **strength** on normal play (? Gemini), but on hard continuous-drilling clips **both** miss ~half (GX010141 recall ~45?52%). The recall lever matters mainly for the hard clips.

45 - ****G3 ? rally-END over-extension**** (quick-return-after-point): shuttle keeps moving ? window runs long / merges; R4 (shuttle-only) can't catch it. Contributes to G2's false positives/merges.

46 - ****Local's gap to Gemini SHRINKS on normal play:**** hard clip ??0.16 (0.51 vs 0.67) ? normal clip ??0.07 (0.82 vs 0.89), and local **out-recalls** Gemini there. The deficit is precision, not detection.

47 - ****What is NOT the gap:**** boundary precision. When local hits a rally its edges are tight (|?start|/?end| ? ~0.8 s, on GX010137 even tighter than Gemini). The R2/R3/R4 windowing is doing its job. [verified]

48 - ****Long-rally lens refinement (P7, n=15, [verified]):**** filtering the eval to >5s rallies (the eye-catching highlights) shows local's single-court over-fire is ****short-FP-dominated**** (seg_ratio 1.43?1.04 ? duration-filterable), but ****multicourt still over-fires 1.58x on LONG rallies**** (adjacent-court games are genuine long rallies). ? the multicourt precision gap needs ****court-localization (Gate-B / court mask, Lever C)****, NOT a duration filter; and the strict tolerance precision is gated by ****over-merge contamination**** of the long-pred set, not just short FPs. See QUALITY_ITERATIONS.md "P7".

49 - ? ****Data caveat:**** GX010141 was labeled before ~3 rally-annotator-tool bugs were found + fixed (GX010137 is post-fix). GX010141's numbers may shift on a re-export ? re-verify before drawing hard conclusions from it.

50

51 ## 4. Design / architecture ? the levers (REUSE MAP)

52 The signal mechanics live in
 MULTI_SIGNAL_FUSION_PLAN.md; this is which lever attacks which gap.

53

54 - **Lever A ? player-state rally-END rule** (targets **G3 + G2**) `[design]`. Generalize R4 from shuttle-only to **shuttle OR bilateral player motion**: end the rally at the last frame with competitive evidence from **either**, and **hard-terminate** on a "shuttle picked up" event. Two cheap detectors, no new model required:

55 - **Player kinematics** ? person velocity drops / players go stationary after a point (the casual returner taps it while the other walks). **Reuse**: the person cue producer (torchvision + ByteTrack) in `yolo_hybrid` (fusion P1, `[verified]` exists).

56 - **Shuttle-in-hand / out-of-play** ? `shuttle_player_colocation` (shuttle box ? person box = held) `[verified]` built in `fusion_features.py` (default-OFF, unmeasured on golden) from the fusion plan §4, **or** the no-new-model proxy: **sustained WASB** `Visibility=0` after the last fast stroke (held shuttle goes dark). Use it as a **termination anchor + reverse-timeline search** to the true end (last net-crossing / last fast-shuttle / last bilateral motion) ? the owner's idea, refined: the hold event is **late**, so anchor-then-backtrack, don't use its timestamp as the end.

57 - **Lever B ? recall** (targets **G1**) `[design]`. Both systems miss real rallies ? detector sensitivity: proposal-windowing / velocity-threshold tuning, and the **complementary-fusion** angle below. Measure recall floor per video; the highest-leverage labels are the **missed** rallies (active learning, RALLY_QUALITY_RESEARCH §4.3).

58 - **Lever C ? signal/ensemble fusion** (targets **G1**) `[design]`. Local & Gemini (and shuttle vs person cues) hit **different** rallies ? a union/ensemble lifts recall. Already-built scaffolding: `rally_gate.py` **Gate A** (net-crossings, structural rally-state, `[verified]` built, eval-only ? wire to prod = cheapest precision win for G2) and the FusionSegmenter feature path.

59 - **Pure-precision cleanup** (**G2**): `rally_gate` Gate A kills single-feed / held-shuttle **starts** structurally (?0?1 net crossings vs a real rally's several) ? the cheapest G2 win and independent of the person stack.

60

61 **Reuse map** (what's new vs there): `rally_gate.py::gate_windows` (Gate A, built) · `tracknet_runner` R4 end-trim (extend, don't replace) · `backend/eval/windowing.py` (preset knobs) · person producer in `yolo_hybrid` (built) · `shuttle_player_colocation` (built in `fusion_features.py`, default-OFF + unmeasured on golden) · `scratch/ingest_golden.py` + `scratch/boxhill17_score_golden.py` (the per-video golden scoring loop, `[verified]`).

62

63 **## 6. Honest limits ? what this does NOT do**

64 - **n is small**. One golden clip scored so far (GX010141); the corpus is n=15 with several weak clips. Per-video conclusions firm up only as the owner labels more (GX010137 / GX020140 / a large-count clip next).

65 - **Nothing here is shipped or even started** ? implementation is owner-gated. A DONE row in §7 means **measured on golden**, not **flipped on by default**.

66 - **Player-state levers need GPU person-detection on the golden clips** (not yet extracted); until then Lever A/B are design, not data.

67 - The boundary tightness is from clips where local **hits** the rally; it says nothing about the rallies it misses (G1).

68

69 **## 7. Deliverables & progress tracker ? source of truth**

70 Legend: ? Todo · ? In progress · ? Done · ? Gated. One small PR per row, off `origin/master`, no stacking.

24. assistant (2026-06-21T13:57:09.599Z)

Let me read the gap-closure §7 tracker (the execution tracker for the remaining gap) to capture the live next-step rows, since this is now the active execution doc peer to RALLY_QUALITY_RESEARCH.

25. user (2026-06-21T13:57:11.124Z)

69 **## 7. Deliverables & progress tracker ? source of truth**

70 Legend: ? Todo · ? In progress · ? Done · ? Gated. One small PR per row, off `origin/master`, no stacking.

71

72 | ID | Deliverable | Depends on | Gated? | Status | PR |

```

73 |----|-----|-----|-----|-----|----|
74 | P0 | Golden scoring loop (`ingest_golden.py` + `boxhill17_score_golden.py`); GX010141
ingested (#12) + scored | ? | No | ? | this work |
75 | P1 | Ingest the rest of the Boxhill golden batch (137/020140 + a large-count clip) ?
per-video baseline table | owner labels | ? owner | ? | ? |
76 | P2 | Wire `rally_gate` **Gate A** (net-crossings) to the served path ? cheapest G2 win
| ? | ? owner | ? | ? |
77 | P3 | Extract person-velocity + `shuttle_player_colocation` on the golden clips
(default-OFF features) | P1 | ? owner | ? | ? |
78 | P4 | **Lever A**: player-state rally-END rule (R-next) ? measure ?F1 / ?merge-rate /
?\\|?end\\| vs R4 | P3 | ? owner | ? | ? |
79 | P5 | **Lever B/C**: recall + complementary fusion ? measure recall lift, no precision
regression | P3 | ? owner | ? | ? |
80 | P6 | Launch decision: milestone + best levers default-ON, with a Launch Report (both
rulers + over-seg guard) | P4,P5 | ? owner | ? | ? |
81 | P7 | **Long-rally (>5s) quality lens** (eval `--min-duration`/`--stratify`) + reel
knob (`stitching.min_rally_duration`), both default-OFF. **Measured n=15:** single-court over-
fire is short-FP-dominated (seg_ratio 1.43?1.04) but **multicourt over-fires 1.58x even on long
rallies** (? court-localization, Lever C) and strict tolP doesn't rise (over-merge-
contaminated). See [`QUALITY_ITERATIONS.md`](QUALITY_ITERATIONS.md) "P7". | ? | No (eval+cfg,
default-OFF) | ? measured | this PR |
82
83 ## 8. Open questions ? owner / external
84 1. **Which lever first** once labels land ? the no-model wins (Gate A wiring + shuttle-
invisibility end-anchor), or go straight to person-velocity? *(decides: owner, after the per-
video baseline)*
85 2. **Per-video target** ? "at-par-with-Gemini-vs-golden on every clip," or an absolute
F1 floor? *(decides: owner)*
86 3. **Recall lever** ? detector-sensitivity tuning vs ensemble fusion as the primary path
for G1. *(decides: data, post-P3)*
87
88 ## 9. Definition of done
89 - [ ] Golden corpus covers the Boxhill set + enough variety that per-video reads are
trustworthy (n meaningfully > 12).
90 - [ ] The **rally-END quick-return** failure mode (G3) is closed: measured ? merge-rate
+ ? |?end| on the clips that exhibit it, no regression elsewhere.
91 - [ ] Local is **at-par-or-better than Gemini vs golden on every corpus video** (no per-
video F1 regression).
92 - [ ] The winning levers are flipped **default-ON** via a Launch Report (over-seg guard
clean on both rulers).
93 - [ ] ? flip `Status ? DONE`, archive to `docs/archives/`.
94
95 ## 10. References
96 **Code anchors `[verified]`**: `backend/pipeline/detectors/rally_gate.py` (Gate A) .
`backend/pipeline/detectors/tracknet_runner.py` (R4 end-trim, `_trim_inactive_tail`) .
`backend/eval/windowing.py` . `backend/eval/rally_seg_eval.py` + `tolerance_metrics.py` (R2) .
`backend/pipeline/segmenters/yolo_hybrid.py` (person cue) . `scratch/ingest_golden.py`,
`scratch/boxhill17_score_golden.py`.
97 **Internal docs**: [`RALLY_QUALITY_RESEARCH.md`](RALLY_QUALITY_RESEARCH.md),
[ `MULTI_SIGNAL_FUSION_PLAN.md` ](MULTI_SIGNAL_FUSION_PLAN.md) §4,
[ `QUALITY_ITERATIONS.md` ](QUALITY_ITERATIONS.md),
[ `EXPERIMENT_HARNESS.md` ](EXPERIMENT_HARNESS.md),
[ `GOLDEN_VIDEOS.md` ](data_pipeline/GOLDEN_VIDEOS.md),
[ `HOW_RALLY_DETECTION_WORKS.md` ](HOW_RALLY_DETECTION_WORKS.md).
98
99 ---
100
101 ### Changelog
102 - 2026-06-18 ? Created. Seeded with the GX010141 golden scoring (first ground-truth
clip), the three named gaps (recall / precision / rally-END), and the player-state end-rule
design from the owner's quick-return insight.
103 - 2026-06-18 ? Added GX010137 (golden #13, normal match play): local F1@0.5 **0.82**
(recall **0.88 > Gemini 0.79**) vs Gemini 0.89. Confirms **opposite failure modes** (local high-
recall/low-precision, Gemini the reverse) ? **precision is local's #1 lever**, fusion the
durable win. Flagged the GX010141 pre-annotator-fix data caveat.

```

104 - 2026-06-18 ? **P7 long-rally (>5s) quality lens + reel knob (both default-OFF),
measured n=15.** Eval `--min-duration`/`--stratify` (`rally\_seg\_eval.py`) + product
`stitching.min\_rally\_duration`. Finding: single-court over-fire is short-FP-dominated (seg_ratio
1.43?1.04) but **multicourt over-fires 1.58× even on long rallies** (? needs court-localization,
Lever C) and strict tolP is gated by over-merge contamination, not just FPs. Sharpened §3 +
added §7 row P7.
105

26. assistant (2026-06-21T13:57:17.735Z)

I now have a complete and current picture of the rally-quality track across all the docs. I have everything needed to produce the structured output. Let me capture the staleness findings I noticed and assemble the answer.

Key findings on freshness: the RALLY_QUALITY_RESEARCH.md §7 tracker is partially stale ? it still shows n=6 and W2/multi-feature scorer as "next (GPU)", while QUALITY_ITERATIONS + R1_MILESTONE_LAUNCH_REPORT + GOLDEN_REGRESSION_FIXTURES show the corpus has grown to n=15, R1 milestone (W1 cap6 + R4) shipped default-ON (PR #204), R2 metric implemented, R5 blob-fallback + P7 long-rally lens shipped default-OFF, and the active execution tracker has moved to RALLY_DETECTION_GAP_CLOSURE.md §7.

27. user (2026-06-21T13:58:31.439Z)

Structured output provided successfully

28. assistant (2026-06-21T13:58:45.847Z)

I've completed the rally-quality track review and returned the structured findings. Summary of what I extracted:

****Status:**** The rally-quality project's windowing wins are shipped. The R1 milestone (W1 cap=6 + R4 rally-END trim) is ****flipped default-ON**** (PR #204, baseline?milestone R2 tolF1 +0.222 at n=13, p=0.0005, over-seg guard PASS, independently reproduced). E1 baseline, W1/W1.5, W2 net-crossings first-step, R2 tolerance metric (augmenting tIoU), R5 blob-fallback, and the P7 long-rally lens are all landed (the latter three default-OFF). The full experiment/ablation harness exists. The active frontier moved to `RALLY\_DETECTION\_GAP\_CLOSURE.md` §7, which re-prioritized the gap to ****detection (precision #1, then recall, then rally-END)**** with ****Gate-A wiring**** as the cheapest precision win.

****Prioritized next steps**** (captured in the structured output): grow golden corpus to n?16 (highest leverage, owner-gated) ? wire Gate-A to prod ? extract person/audio features on golden clips ? W2 multi-feature learned scorer LOVO ? R2-vs-tIoU ablation to promote R2 as the gate + IAA ?-calibration ? Lever A (player-state rally-END rule) ? Lever B/C fusion ? M2 per-venue calibration ? finish Golden Regression Fixtures M4/P0. Tier-3 owned boundary-head is a corpus-gated stretch.

****Key gates:**** most ML steps are ****blocked on GPU-extracted golden features + labels (owner-side, not in this checkout)****; the gap-closure rows P1-P6 are ****owner-approval-gated****; the ship-gate target (#172) is still an open owner decision.

****Staleness flagged (not fixed):**** `RALLY\_QUALITY\_RESEARCH.md` §7 tracker still says n=6 and lists W2 as "next" while the corpus is n=15 and R1/R2/R5/P7 have shipped ? it should be reconciled or point to the now-live `RALLY\_DETECTION\_GAP\_CLOSURE.md` §7. The `NEXT\_STEPS.md` rally-quality detail body is dated 2026-06-15 (pre-milestone). Corpus `n` is cited inconsistently (6/13/15) across docs, and the headline gap-closure % is still uncomputed (Gemini not measured on the same golden set).

One process note: I could not `git pull --ff-only` ? it failed with "Cannot fast-forward to multiple branches," so I worked read-only against the local checkout; the docs may lag the remote.