

claude-windows / c7679509-ee2b-4443-b240-0b3b1b1a7291

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\subagents\workflows\wf_a2b97cff-dcd\agent-af065cf7f829b31bd.jsonl`  
- SHA-256: `ec4ea28f394bd316653c6806ea84fae8d885185669c24015273d1d20a3762094`  
- Source modified: `2026-06-21T14:00:54+00:00`  
- Imported at: `2026-07-05T16:48:23+00:00`  
- project: `wf_a2b97cff-dcd`  
- session_id: `c7679509-ee2b-4443-b240-0b3b1b1a7291`
```

Transcript

1. user (2026-06-21T13:56:47.022Z)

Adversarially VERIFY the load-bearing claims below about packaging/distributing our Python app. For EACH claim, try to REFUTE it (check feasibility, licensing accuracy, version facts, whether it actually applies to a torch+opencv+ffmpeg+FastAPI app). Default to "uncertain" if you cannot confirm. Use web search to check current facts.

TOPIC: dual-python-env

CLAIMS:

```
[  
  {  
    "claim": "WASB-SBDT pins exactly numpy==1.22.4, torch==1.11.0+cu113, torchvision==0.12.0+cu113 (cu113 extra-index), plus hydra-core==1.2.0, timm==0.6.5, einops==0.4.1, opencv-python==4.6.0.66 ? i.e. the second env is genuinely a torch-1.11/CUDA-11.3-era stack, distinct from the engine's modern torch.",  
    "confidence": "high",  
    "source": "https://github.com/nttcom/WASB-SBDT/blob/main/Dockerfile (read via gh api: `pip3 install numpy==1.22.4 torch==1.11.0+cu113 torchvision==0.12.0+cu113`)"  
  },  
  {  
    "claim": "WASB-SBDT is MIT licensed (Copyright 2023 NTT Communications) ? code is commercial-clean under the MIT/Apache/BSD guardrail.",  
    "confidence": "high",  
    "source": "https://github.com/nttcom/WASB-SBDT/blob/main/LICENSE.md"  
  },  
  {  
    "claim": "The WASB model is a standard CNN (HRNet-style backbone via timm + an online tracker), no exotic ops ? src/ has models/, detectors/, trackers/, dataloaders/; timm==0.6.5 is the backbone source. This is what makes both torch-2.x port (d) and ONNX export (d') low-op-risk.",  
    "confidence": "high",  
    "source": "github gh api repos/nttcom/WASB-SBDT/contents/src (configs, detectors, models, trackers) + Dockerfile timm==0.6.5"  
  },  
  {  
    "claim": "The Anaconda ToS trap is the `defaults` channel (repo.anaconda.com), which requires a paid license for orgs with 200+ employees ? NOT conda the tool. conda-forge / miniforge / micromamba are BSD-3-Clause and free for ALL use including commercial redistribution.",  
    "confidence": "high",  
    "source": "https://github.com/conda-forge/miniforge/blob/main/LICENSE + https://pydevtools.com/handbook/explanation/is-conda-actually-free/ + https://ubinfie.github.io/2024/10/15/anaconda-defaults.html"  
  },  
  {  
    "claim": "conda-pack is New BSD licensed and creates relocatable conda envs, but the build OS must match the target OS and an env cannot be relocated again after conda-unpack.",
```

```

    "confidence": "high",
    "source": "https://conda.github.io/conda-pack/ + https://github.com/conda/conda-pack"
  },
  {
    "claim": "micromamba is a fully static, self-contained ~single-binary env manager (BSD)
ideal for bundling; it can create envs from a conda-forge lock file with no conda install
present.",
    "confidence": "high",
    "source": "https://mamba.readthedocs.io/en/stable/user_guide/micromamba.html +
https://github.com/mamba-org/mamba"
  },
  {
    "claim": "Bundling CUDA torch is ~4?5 GB on disk per env (torch GPU wheel ~583 MB, grows to
~4 GB bundled with CUDA/cuDNN libs) ? so two coexisting torch envs is roughly a 10 GB footprint;
this is the core reason to prefer collapsing to one env or to ONNX.",
    "confidence": "high",
    "source": "https://github.com/orgs/pyinstaller/discussions/8552 +
https://github.com/pytorch/pytorch/issues/17621"
  },
  {
    "claim": "PyTorch officially guarantees state_dict backward-compat: a state_dict saved on
1.11 loads into a 2.x model built with the same architecture; standard CNN ops (conv/pool/relu)
have no breaking changes across versions. (Loading a whole pickled model is NOT guaranteed ? use
state_dict.)",
    "confidence": "high",
    "source": "https://docs.pytorch.org/docs/2.11/generated/torch.nn.Module.html (backward-
compat note) + pytorch.org/get-started/pytorch-2-x"
  },
  {
    "claim": "torch.onnx.export produces a self-contained graph that runs with onnxruntime
WITHOUT the torch framework; onnxruntime offers CPU, CUDA, DirectML (Windows GPU) and CoreML
execution providers and is MIT-licensed ? a small (tens-of-MB) cross-platform runtime.",
    "confidence": "high",
    "source": "https://onnxruntime.ai/blogs/pytorch-on-the-edge +
https://onnxruntime.ai/docs/tutorials/export-pytorch-model.html"
  },
  {
    "claim": "The engine ALREADY ships the subprocess seam (c) and the ABC explicitly documents
an ONNXRunner (d') de-risk path: native_wasb_runner.py invokes WASB_PYTHON by path and shells
wasb_infer.py to a Frame,Visibility,X,Y CSV; base.py:46-54 names a 'LinuxNativeRunner (or
ONNXRunner)' that loads weights 'torch or onnx' and writes 'the exact same CSV format so
trajectory_to_action_windows works verbatim'. Owner byte-parity validated 2026-06-20 (1194/1195
frames).",
    "confidence": "high",
    "source": "C:/Users/avidu/Projects/badminton-highlight-
indexer/backend/pipeline/detectors/native_wasb_runner.py:52,110-113,210-237 + base.py:46-54"
  },
  {
    "claim": "The engine's main env is modern torch 2.x (requirements.txt pins torch>=2.12.0 /
torchvision>=0.27.0; pyproject keeps torch OUT of deps because CUDA wheels need --index-url) ?
confirming the version clash with WASB's 1.11 is real and unavoidable in a single shared
interpreter.",
    "confidence": "high",
    "source": "C:/Users/avidu/Projects/badminton-highlight-indexer/requirements.txt
(torch>=2.12.0) + pyproject.toml:8-11,58-65"
  },
  {
    "claim": "Delivery is a self-hosted webserver+UI the user spins up (NOT a desktop app),
targeting non-tech recreational players on commodity (often no-GPU) hardware; D3 already commits
to dropping WSL and running the detector natively per-OS; D4 requires MIT/Apache/BSD-only
bundled deps incl. weights.",
    "confidence": "high",
    "source": "C:/Users/avidu/Projects/khelsutra/docs/LOCAL_APPLIANCE_ARCHITECTURE.md D10 +
C:/Users/avidu/Projects/badminton-highlight-indexer/docs/DESKTOP_APP_PLAN.md D3,D4 §3.1"
  }
}

```

]

RECOMMENDATION UNDER TEST: Layer two decisions, do NOT pick a single bullet.

1) PACKAGING TIERING = (f). Ship a torch-FREE LIGHT tier as the default download: google-genai is already a hard dep, so Gemini + motion + stub deliver the full ingest -> rallies -> reel CUJ with zero torch, a tiny install, and a reliable download for the non-tech India-coach persona on a slow uplink. The on-device WASB detector is an OPT-IN HEAVY add-on. This preserves every completed CUJ and matches the appliance's already-shipped 'runs on Gemini today' reality (LOCAL_APPLIANCE_ARCHITECTURE D10).

2) For the HEAVY tier, the cleanest long-term answer is (d') EXPORT WASB TO ONNX + onnxruntime ? which the codebase was explicitly designed for (base.py names the ONNXRunner; native_wasb_runner.py is the sibling that already shells out to the same CSV contract). It DISSOLVES the dual-env problem entirely: the WASB path needs neither torch 1.11 nor torch 2.x, ships a ~tens-of-MB runtime instead of a 5 GB framework, and gives CPU + NVIDIA(CUDA) + Windows(DirectML) + Apple(CoreML) from one artifact ? directly answering the no-GPU-laptop and Apple-Silicon unknowns the DESKTOP_APP_PLAN P0 spike was created to resolve.

3) BRIDGE / fallback while ONNX export is validated: collapse to one env via (d) torch-2.x port if the timm/HRNet load is quick, OR ship the second env via (b) micromamba-from-conda-forge (license-clean, sidesteps the Anaconda `defaults` ToS) reusing the EXISTING subprocess seam (c) unchanged ? pointing WASB_PYTHON at the micromamba env. Use conda-pack (a) only as a build-time convenience for an offline-friendly first-run (no network), but mind its OS-lock + no-relocate caveat. AVOID (e) Docker as the LOCAL default ? it requires Docker + NVIDIA toolkit, a non-starter for the target user (keep Docker only for the CLOUD/HEAVY worker, where it aligns with the M6 Cloud Run + GPU seam).

LICENSING NON-NEGOTIABLE: whatever env path is chosen, build/solve from conda-forge ONLY (never the Anaconda `defaults` channel) ? micromamba/miniforge/conda-pack tools are all BSD; the ToS trap is the channel, not the tool. WASB-SBDT (MIT), timm (Apache), onnxruntime (MIT) are all clean.

Both ONNX (d') and torch-port (d) MUST re-pass the owner byte-parity / golden-regression gate on a real GPU before production ? numeric drift across torch versions / the ONNX op set is the one real risk, and the offline mocked suite cannot prove the numbers (hand to owner with the labelled corpus).

2. assistant (2026-06-21T13:56:50.092Z)

I'll verify each claim adversarially. Let me start by examining the local codebase claims (which I can verify directly) and then web-check the external facts.

3. user (2026-06-21T13:56:51.060Z)

```
1      """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3      The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4      "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
M3.5).
5      It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
6      Linux GPU box (GCP / a SaaS worker): it runs the same verified ``wasb_infer.py``
core
7      that the WSL adapter runs, so the trajectory CSV is identical (to a sub-pixel float
tolerance;
8      see the acceptance gate below) ? but WITHOUT the WSL transport:
9
10     * NO ``wslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
11     ``wasb`` env's python is invoked by path), NO ``/mnt/c`` path translation
12     (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
13     * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps``, ``cpu``, or
``auto``
14     (M4: cuda-if-available-else-cpu ? the CPU fallback for a no-GPU box; explicit
``cuda`` stays
15     strict); threaded into ``wasb_infer.py --device`` (and so the Hydra
``runner.device`` override).
```

```

16     * weights / repo / python / scratch come from the **environment** first
17     (`WASB_WEIGHTS` / `WASB_REPO_DIR` / `WASB_PYTHON` / `WASB_SCRATCH`), then
18     `indexing.wasb.*` config, then a default ? so a box needs no config edits.
19
20     It emits the identical `Frame,Visibility,X,Y` CSV (same `{stem}__{vid12}_wasb.csv`
21     name as the WSL runner) and reuses the model-agnostic `parse_trajectory_csv` /
22     `trajectory_to_action_windows` from `tracknet_runner`, so the trajectory hybrids and
23     the windowing brain are unchanged. Selected via the existing `detector_impl` seam
24     (`trajectory_hybrid._resolve_runner` ? `_build_native_runner`);
25     `detector_impl="native"`.
26
27     **Acceptance gate (parity, NOT strict byte-equality):** the SAME clip through the WSL
28     runner
29     (Windows) and this runner (Linux, `device=cuda`) yields the same trajectory CSV + the
30     same
31     candidate windows **within a small float tolerance (~1e-3 px)**. cuDNN is non-
32     deterministic
33     ACROSS GPUs, so a sub-pixel diff is expected on different hardware ? but it is byte-
34     exact
35     run-to-run on a FIXED GPU. ? Validated 2026-06-20 on an RTX 2070 (via WSL): native
36     streaming
37     vs the cached WSL-disk CSV matched **1194/1195** frames (the lone diff a ~5e-5 px edge
38     artifact), and two native runs were byte-identical (deterministic). This is a hardware
39     check
40     (a real GPU + the `wasb` env), so the offline, subprocess-mocked suite asserts the
41     contract,
42     not the numbers.
43     """
44     import os
45     import shutil
46     import logging
47     import subprocess
48     from dataclasses import dataclass
49     from typing import Any, Dict, List, Optional, Tuple
50
51     from backend.pipeline.detectors.base import DetectorRunner
52     from backend.pipeline.detectors.device import AUTO, VALID_DEVICES, DeviceContext
53     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
54     WSL/TrackNet.
55     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
56
57     logger = logging.getLogger(__name__)
58
59     # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
60     # detectors/trackers/data loaders + find configs/), exactly as the WSL adapter does.
61     _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
62
63     # Single torch/cuda probe code, used by BOTH healthcheck and the lazy device-resolution
64     probe
65     # so the "cuda True/False" string they parse can never drift apart.
66     _CUDA_PROBE_CODE = "import torch; print('torch', torch.__version__, 'cuda',
67     torch.cuda.is_available())"
68
69     @dataclass
70     class NativeWasbConfig:
71         """Resolved settings for a Linux-native WASB run.
72
73         Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
74         over `indexing.wasb.*` config (the box sets env; no config edits needed).
75         `device`
76         defaults to `cuda` (the WSL parity path); `stream_video` defaults to True (the
77         perf win).
78         """
79         repo_dir: str = "~/models/WASB-SBDT" # native clone of nttcom/WASB-SBDT

```

```

68         python_bin: str = "python"                # the `wasb` conda env python (invoked by
path, no activate)
69         weights_path: str = ""                    # REQUIRED (env WASB_WEIGHTS or
indexing.wasb.weights_path)
70         sport: str = "badminton"
71         device: str = "cuda"                      # auto | cuda | mps | cpu (auto = cuda-if-
available-else-cpu)
72         scratch_dir: str = ""                     # frame-extraction scratch (env); "" = next
to the output dir
73         timeout_sec: int = 1800                   # 30 min; a hung GPU call is killed (0 = no
timeout)
74         keep_frames: bool = False                 # retain the decoded frame cache after
success (debug only)
75         # NATIVE DEFAULT = STREAMING: skip the cv2 PNG extraction that dominates a long
clip's
76         # wall-clock (measured ~50s on a 20s clip) ? bit-identical to disk (verified on a
real GPU
77         # 2026-06-20). The GPU box reads from a local NVMe so streaming is the right default
here.
78         stream_video: bool = True
79
80     @classmethod
81     def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "NativeWasbConfig":
82         w = (idx_cfg or {}).get("wasb", {}) or {}
83         env = os.environ.get
84         # stream_video is tri-state in config: None/absent => native default (stream);
an
85         # explicit True/False still wins (e.g. False for a container cv2 can't seek).
86         sv = w.get("stream_video")
87         return cls(
88             repo_dir=env("WASB_REPO_DIR") or w.get("repo_dir", "~/models/WASB-SBDT"),
89             python_bin=env("WASB_PYTHON") or w.get("python_bin", "python"),
90             weights_path=env("WASB_WEIGHTS") or w.get("weights_path", "") or "",
91             sport=w.get("sport", "badminton"),
92             device=env("WASB_DEVICE") or w.get("device", "cuda"),
93             scratch_dir=env("WASB_SCRATCH") or env("RALLY_SCRATCH") or env("TMPDIR") or
"",
94             timeout_sec=int(w.get("timeout_sec", 1800)),
95             keep_frames=bool(w.get("keep_frames", False)),
96             stream_video=(True if sv is None else bool(sv)),
97         )
98
99
100     class NativeWasbRunner(DetectorRunner):
101         """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
docstring."""
102
103         def __init__(self, cfg: NativeWasbConfig):
104             self.cfg = cfg
105             # Cached CUDA-availability probe (used to resolve device="auto"). Set by
healthcheck;
106             # lazily probed by run_predict if healthcheck was skipped. None = not yet
probed.
107             self._cuda_available: Optional[bool] = None
108
109             # --- low-level native exec (the single place tests patch)
-----
110             def _run(self, argv: List[str], cwd: Optional[str] = None) ->
subprocess.CompletedProcess:
111                 logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
112                 return subprocess.run(argv, cwd=cwd, capture_output=True, text=True,
113                                     timeout=(self.cfg.timeout_sec or None))
114
115             # --- device resolution (M4 DeviceContext: device="auto" ? cuda-if-available-else-
cpu) -----

```

```

116     def _probe_cuda(self) -> bool:
117         """Probe ``torch.cuda.is_available()`` in the wasb env's python (a subprocess).
Any failure
118         (missing python / timeout / import error) is treated as 'no CUDA'."""
119         try:
120             res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
121         except (FileNotFoundError, subprocess.TimeoutExpired):
122             return False
123         return res.returncode == 0 and "cuda True" in (res.stdout or "")
124
125     def _cuda_is_available(self) -> bool:
126         """CUDA availability, cached across healthcheck + run_predict so we probe at
most once."""
127         if self._cuda_available is None:
128             self._cuda_available = self._probe_cuda()
129         return self._cuda_available
130
131     def _effective_device(self) -> str:
132         """The torch device to actually pass to ``wasb_infer.py``. Only ``auto`` needs
the CUDA
133         probe; an explicit cuda/mps/cpu resolves to itself (so ``device='cuda'`` is
unchanged)."""
134         if self.cfg.device == AUTO:
135             return DeviceContext(AUTO, self._cuda_is_available()).effective
136         return self.cfg.device
137
138     def _repo_src(self) -> str:
139         return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
140
141     def _copy_infer_script(self) -> bool:
142         """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
configs/."""
143         repo_src = self._repo_src()
144         try:
145             os.makedirs(repo_src, exist_ok=True)
146             shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
147             return True
148         except OSError as e:
149             logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
150         return False
151
152     def _rm_rf(self, path: Optional[str]) -> None:
153         """Best-effort recursive delete of a native path (post-success cleanup; never
raises)."""
154         if path:
155             shutil.rmtree(path, ignore_errors=True)
156
157     def healthcheck(self) -> Tuple[bool, str]:
158         """Verify weights + repo present and the `wasb` python imports torch (and sees a
GPU on cuda)."""
159         if not self.cfg.weights_path:
160             return False, ("WASB weights path is empty ? set WASB_WEIGHTS (or
indexing.wasb.weights_path).")
161         weights = os.path.expanduser(self.cfg.weights_path)
162         if not os.path.exists(weights):
163             return False, f"WASB weights not found at {weights}"
164         repo_src = self._repo_src()
165         if not os.path.isdir(repo_src):
166             return False, (f"WASB-SBDT repo src/ not found at {repo_src} ? clone
nttcom/WASB-SBDT "
167                             f"(set WASB_REPO_DIR / indexing.wasb.repo_dir).")
168         # Fail fast on a typo'd device (the device-policy seam) ? a clear error here
beats a cryptic
169         # argparse rc=2 from wasb_infer.py at run time.
170         if self.cfg.device not in VALID_DEVICES:

```

```

171         return False, (f"unknown device {self.cfg.device!r} ? valid: {'',
172         '.join(VALID_DEVICES)).")
173     try:
174         res = self._run([self.cfg.python_bin, "-c", _CUDA_PROBE_CODE])
175     except FileNotFoundError:
176         return False, f"python binary not found: {self.cfg.python_bin!r} (set
177         WASB_PYTHON)."
178     except subprocess.TimeoutExpired:
179         return False, "native torch healthcheck timed out."
180     out = (res.stdout or "").strip()
181     if res.returncode != 0:
182         return False, f"`{self.cfg.python_bin}` torch import failed: {(res.stderr or
183         out).strip()}"
184     # M4: resolve the device. Cache the probe so run_predict reuses it (no second
185     subprocess).
186     self._cuda_available = "cuda True" in out
187     ctx = DeviceContext(self.cfg.device, self._cuda_available)
188     if ctx.cuda_required_but_missing:
189         # Unchanged pre-M4 behaviour for an EXPLICIT cuda request: hard-fail, never
190         a silent
191         # slow-CPU downgrade. Set indexing.wasb.device=auto for a CPU fallback.
192         return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
193     if ctx.fell_back:
194         logger.warning("device=auto: no CUDA GPU detected (%s) ? FALLING BACK TO
195         CPU; WASB "
196         "inference will be slow. Set indexing.wasb.device=cuda to
197         require a GPU.", out)
198     return True, out
199
200 def run_predict(self, video_win_path: str, output_win_dir: str,
201                 video_id: Optional[str] = None) -> Optional[str]:
202     """Run WASB natively on a video. Returns the path to the trajectory CSV, or
203     None.
204     Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
205     WSL
206     runner does, and the CSV is named ``{stem}__{vid12}_wasb.csv`` ? byte-for-byte
207     the
208     same downstream contract, so the parity comparison is apples-to-apples.
209     """
210     output_dir = os.path.abspath(output_win_dir)
211     os.makedirs(output_dir, exist_ok=True)
212     # Normalize for a cross-platform basename (tests pass Windows paths on Linux).
213     norm = str(video_win_path).replace("\\", "/")
214     stem, _ext = os.path.splitext(os.path.basename(norm))
215     # Content-derived key: same filename + different bytes -> different key (no
216     cache reuse).
217     key = f"{stem}__{video_id[:12]}" if video_id else (stem or "wasb")
218     out_csv = os.path.join(output_dir, f"{key}_wasb.csv")
219
220     if not self._copy_infer_script():
221         logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
222         return None
223
224     weights = os.path.expanduser(self.cfg.weights_path)
225     argv = [self.cfg.python_bin, "wasb_infer.py",
226            "--video", os.path.abspath(str(video_win_path))]
227     frames_dir: Optional[str] = None
228     if self.cfg.stream_video:
229         # Fast path: decode in-process, no PNG extraction. Output is bit-identical
230         to disk.
231         argv.append("--stream-video")
232     else:
233         scratch = self.cfg.scratch_dir or output_dir
234         frames_dir = os.path.join(scratch, f"{key}_frames")

```

```

224         argv += ["--frames_out_dir", frames_dir]
225         # M4: resolve device="auto" to the effective device (cuda-if-available-else-
cpu); an
226         # explicit cuda/mps/cpu passes through unchanged. Never sends "auto" to
wasb_infer.py.
227         device = self._effective_device()
228         argv += ["--weights", weights,
229                 "--sport", self.cfg.sport,
230                 "--device", device,
231                 "--out", out_csv]
232
233         logger.info("Running native WASB inference (device=%s%s) on %s ...",
234                     device, " [auto?fallback]" if self.cfg.device == AUTO and device !=
"cuda" else "",
235                     os.path.basename(norm))
236         try:
237             res = self._run(argv, cwd=self._repo_src())
238         except FileNotFoundError:
239             logger.error("python binary not found: %r (set WASB_PYTHON).",
self.cfg.python_bin)
240             return None
241         except subprocess.TimeoutExpired:
242             logger.error("WASB native inference timed out after %ss.",
self.cfg.timeout_sec)
243             return None
244         if res.returncode != 0:
245             logger.error("WASB native inference failed:\n%s", (res.stderr or res.stdout
or "")[-2000:])
246             return None
247         if not os.path.exists(out_csv):
248             logger.error("WASB native finished but expected CSV not found at %s",
out_csv)
249             return None
250         logger.info("WASB trajectory CSV: %s", out_csv)
251
252         # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
253         # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
254         # (Streaming never materializes frames, so there is nothing to drop on that
path.)
255         if not self.cfg.keep_frames and frames_dir:
256             logger.info("Cache hygiene: removing frame cache for %s "
257                         "(set indexing.wasb.keep_frames=true to retain).", key)
258             self._rm_rf(frames_dir)
259         return out_csv
260
261         # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
262         parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
263         trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)
264

```

4. user (2026-06-21T13:56:52.050Z)

```

1     fastapi>=0.111.0
2     uvicorn>=0.30.1
3     opencv-python>=4.9.0.80
4     numpy>=1.26.4
5     pydantic>=2.7.1
6     jinja2>=3.1.4
7     python-dotenv>=1.0.0
8     google-genai>=2.4.0
9     torch>=2.12.0

```



```

10     torchvision>=0.27.0
11     supervision>=0.21.0,<0.30.0
12
13     # Per-video observability reports (SOFT dependency: the pipeline degrades to a
14     # no-op if absent). Shared, permissively-licensed report toolkit. Pin a tag once
15     # the private repo is published, e.g.:
16     #   obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3
17     # For local development install it editable from a sibling checkout:
18     #   pip install -e ../sports-obsreport --no-build-isolation
19

```

5. user (2026-06-21T13:56:52.142Z)

```

1     # PEP 621 packaging for the badminton/sports highlight indexer (DEFERRED #7).
2     #
3     # Build backend: hatchling. `backend/` is an intentional NAMESPACE package
4     # (no `backend/__init__.py`; mypy.ini sets explicit_package_bases for the same
5     # reason), so the wheel target packages are listed EXPLICITLY below ? hatchling's
6     # auto-detection would otherwise miss the un-__init__'d `backend` root.
7     #
8     # torch/torchvision are deliberately NOT listed in [project.dependencies]:
9     # their CPU/CUDA wheels need an out-of-band index (--index-url / --extra-index-url),
10    # which PEP 621 cannot express. Install them separately ? see [project.optional-
dependencies]
11    # `gpu` (documentation only) and the README. CI installs CPU torch on its own line.
12    [build-system]
13    requires = ["hatchling"]
14    build-backend = "hatchling.build"
15
16    [project]
17    name = "badminton-highlight-indexer"
18    version = "0.1.0"
19    description = "AI-assisted rally segmentation and highlight indexing for racket sports."
20    readme = "README.md"
21    requires-python = ">=3.10"
22    license = { text = "MIT" }
23    authors = [{ name = "Avi Dullu", email = "avi.dullu@gmail.com" }]
24    keywords = ["badminton", "sports", "video", "highlights", "rally", "segmentation"]
25
26    # Runtime deps mirror requirements.txt, EXCEPT torch/torchvision (see header).
27    dependencies = [
28        "fastapi>=0.111.0",
29        "uvicorn>=0.30.1",
30        "opencv-python>=4.9.0.80",
31        "numpy>=1.26.4",
32        "pydantic>=2.7.1",
33        "jinja2>=3.1.4",
34        "python-dotenv>=1.0.0",
35        "google-genai>=2.4.0",
36        "supervision>=0.21.0,<0.30.0",
37    ]
38
39    [project.optional-dependencies]
40    # Test + lint/type tooling ? matches the CI lanes (.github/workflows/ci.yml).
41    dev = [
42        "pytest",
43        "pytest-cov",
44        "httpx",          # fastapi TestClient transport
45        "ruff",
46        "mypy",
47    ]
48
49    # Per-video observability reports (SOFT dependency: the pipeline degrades to a
50    # no-op if absent). Pinned to its git tag per BACKLOG. The private repo is
51    # fetched over SSH; if you lack access, simply omit this group ? reporting

```

```

52 # tests skip via pytest.importorskip. For local dev against a sibling checkout:
53 # pip install -e ../sports-obsreport --no-build-isolation
54 reporting = [
55     "obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3",
56 ]
57
58 # `gpu` is DOCUMENTATION ONLY ? it is intentionally empty. torch/torchvision
59 # CUDA wheels require `--index-url`/`--extra-index-url`, which PEP 621 cannot
60 # express, so they are NOT declarable here. Install them out-of-band instead:
61 # # CUDA 12.4 box:
62 # pip install torch>=2.12.0 torchvision>=0.27.0 --index-url
https://download.pytorch.org/whl/cu124
63 # # CPU-only box (also what CI uses):
64 # pip install torch torchvision --index-url https://download.pytorch.org/whl/cpu
65 gpu = []
66
67 # Storage backends (docs/DECOUPLED_COMPUTE/). The default `local` backend needs NONE of
68 # these. Cloud SDKs are imported lazily, so `import backend.storage` works without them;
69 # install the extra for the bucket you use. (boto3's S3 client covers AWS + Cloudflare
R2 +
70 # DO Spaces + MinIO via endpoint_url.) `gdrive://` needs NO extra ? it reads a locally
71 # MOUNTED Google Drive (Google Drive for Desktop) as plain files; see
backend/storage/gdrive.py.
72 storage-s3 = ["boto3>=1.34"]
73 storage-gcs = ["google-cloud-storage>=2.16"]
74
75 [project.scripts]
76 # `indexer --server`, `indexer --help`, etc. -> backend/main.py::main()
77 indexer = "backend.main:main"
78
79 [tool.hatch.metadata]
80 # The `reporting` extra pins obsreport to a git tag (a PEP 508 *direct
81 # reference*); hatchling requires opting in to allow those in metadata.
82 allow-direct-references = true
83
84 [tool.hatch.build.targets.wheel]
85 # `backend` is a namespace package (no __init__.py); list both import roots
86 # explicitly so an editable/wheel install exposes `import backend` and
87 # `import training` regardless of cwd.
88 packages = ["backend", "training"]
89

```

6. assistant (2026-06-21T13:57:01.124Z)

Let me verify the `base.py` ABC claims and check the WASB runner companion file, plus the cited docs.

7. user (2026-06-21T13:57:02.471Z)

```

1     """Abstract base for detector runners (TrackNet, WASB, future native, etc.).
2
3     This completes the Detector Runner Abstraction (low-regret item 6 continuation / item 1
4     of the post-observability re-prioritization) to de-risk the WSL/conda dependency for
5     the high-accuracy shuttle trajectory substrate used by tracknet_hybrid / wasb_hybrid.
6
7     The ABC lets us:
8     - Keep the current WSL adapters working unchanged for Windows users.
9     - Swap in (or A/B test) a Linux-native implementation later with zero changes to
10    the hybrid segmenters or the model-agnostic windowing brain.
11    - Provide a single place to evolve the healthcheck / error contract and thread
12    detector readiness into per-video reports.
13
14    Contract (deliberately matches what the concrete runners + hybrids + tests already did):
15    - run_predict(video_win_path: str, output_win_dir: str) -> Optional[str]
16    Run detector, return the Windows-side path to the produced trajectory CSV

```

```

17     (e.g. ".../clip_predict.csv" or ".../clip_wasb.csv"), or None on any failure.
18     The caller (hybrid) is responsible for output dir; impls may stage the video
19     into WSL fs for perf.
20
21     - healthcheck() -> tuple[bool, str]
22     Return (ok, message). Never raises. "ok" decides whether to proceed to run_predict.
23     Message is logged (and can be surfaced to reports or UX in future).
24     Intended to be cheap (env existence, import, GPU visibility, weights present).
25
26 Existing behavior is 100% preserved. The only additions are the inheritance and
27 the explicit common type for future implementers.
28
29 How to add a new runner (for docs / future contributors):
30 1. Subclass DetectorRunner in a new or existing module under pipeline/detectors/.
31 2. Implement run_predict and healthcheck (return (bool, str) tuple).
32 3. If it produces the same CSV schema (Frame,Visibility,X,Y), reuse or re-export
33    TrackNetRunner.parse_trajectory_csv and .trajectory_to_action_windows (they are
34    deliberately static and model-agnostic). If the new detector needs different
35    windowing, it can implement its own or we generalize later.
36 4. Lazy-import heavy native deps inside the methods (so importing the segmenter
37    registry never pulls torch/tf on a machine that doesn't have them).
38 5. Update the two hybrid segmenters (or introduce a detector= choice) to accept
39    DetectorRunner instances (they already do the right thing via duck-typing + the
40    type hints we add here).
41 6. Add unit tests that exercise the contract (isinstance, healthcheck shape, run
42    returns str|None) using mocks for the external bits ? see test_tracknet_runner.py.
43 7. No change to eval/calibration code paths (they consume the CSV + the static
44    windowing helpers directly).
45
46 Optional Linux-native path (the main de-risk goal):
47 A LinuxNativeRunner (or ONNXRunner) would:
48 - Skip all wsl / _stage / conda activate.
49 - Load weights locally (torch or onnx) or call a native binary.
50 - Write the exact same CSV format so trajectory_to_action_windows works verbatim.
51 - healthcheck can just check "weights file exists and torch importable + CUDA?".
52 This removes the WSL requirement for Linux users / future SaaS workers while
53 keeping the exact same candidate quality characteristics (or better, once native
54 weights are available).
55
56 Error surfacing into reports:
57 The healthcheck message is already logged by the hybrids on failure. The reporting
58 harvest sees the outcome as "0 candidates" + detector name. When we wire richer
59 segmenter Results (or pass detector_health into emit_indexer_report context), the
60 (ok, msg) can be recorded explicitly under segmentation.details or a new detector
61 stage. The ABC makes that future change local to one place.
62
63 See also:
64 - backend/pipeline/segmenters/{tracknet_hybrid,wasb_hybrid}.py (the two callers)
65 - tests/test_tracknet_runner.py (contract + pure logic tests; hybrids mock the runner)
66 - docs/TRACKNET_WSL_SETUP.md and CODE_MAP.md (2.2 Detectors section)
67 """
68
69 import logging
70 import shlex
71 import subprocess
72 from abc import ABC, abstractmethod
73 from typing import Any, Callable, List, Optional, Tuple
74
75 logger = logging.getLogger(__name__)
76
77
78 def wsl_tilde_quote(path: str) -> str:
79     """shlex.quote that still lets bash expand a leading `~/`.
80
81     ``shlex.quote("~/clips/x.mp4")`` single-quotes the whole path, so bash

```

```

82     treats ``~`` literally and ``cp``/``mkdir`` fail with "No such file or
83     directory" ? this broke live WSL staging for the default ``~/clips``
84     stage dir (found by the 2026-06-11 wasb_hybrid smoke run; the BACKLOG
85     "shlex-quote breaks ~ expansion" item). Quoting only the part after
86     ``~/`` keeps expansion AND protects spaces/specials in the remainder.
87     """
88     if path == "~":
89         return path
90     if path.startswith("~/"):
91         return "~/\" + shlex.quote(path[2:])
92     return shlex.quote(path)
93
94
95 class WslCommandMixin:
96     """Shared ``wsl -d <distro> bash -c 'source <conda_sh> && <cmd>'`` invocation.
97
98     Lives beside (not on) DetectorRunner: a future Linux-native runner must not
99     inherit WSL plumbing. Requires ``self.cfg`` with ``distro``, ``conda_sh``
100    and ``timeout_sec`` attributes (TrackNetConfig and WasbConfig both qualify).
101    ``timeout_sec == 0`` maps to None = wait forever (configs default 1800s).
102    """
103
104    cfg: Any
105
106    def _wsl_bash(self, inner_cmd: str) -> subprocess.CompletedProcess:
107        full = f"source {self.cfg.conda_sh} && {inner_cmd}"
108        argv = ["wsl", "-d", self.cfg.distro, "bash", "-c", full]
109        logger.debug("WSL exec: %s", full)
110        return subprocess.run(argv, capture_output=True, text=True,
111                               timeout=(self.cfg.timeout_sec or None))
112
113    # ---- cache hygiene helpers (shared by both WSL runners) -----
114    def _wsl_free_gb(self, wsl_path: str) -> Optional[float]:
115        """Best-effort free space (GB) on the WSL filesystem holding ``wsl_path``.
116
117        Returns None if it can't be determined (never raises). Used as a pre-run
118        disk guard so a long extraction doesn't fill the disk mid-flight.
119        """
120        try:
121            res = self._wsl_bash(f"df -P -B1 {wsl_tilde_quote(wsl_path)} | awk
122            'NR==2{{print $4}}'")
123            except (subprocess.TimeoutExpired, OSError):
124                return None
125            try:
126                return int((res.stdout or "").strip()) / (1024 ** 3)
127            except (ValueError, AttributeError):
128                return None
129
130    def _wsl_rm_rf(self, wsl_paths: List[str]) -> None:
131        """Best-effort recursive delete of WSL paths (post-success cleanup; never
132        raises)."""
133        paths = [p for p in wsl_paths if p]
134        if not paths:
135            return
136        quoted = " ".join(wsl_tilde_quote(p) for p in paths)
137        try:
138            self._wsl_bash(f"rm -rf {quoted}")
139        except (subprocess.TimeoutExpired, OSError) as e:
140            logger.warning("WSL cache cleanup failed (non-fatal): %s", e)
141
142    def wsl_source_unreadable_reason(self, src_mnt: str) -> Optional[str]:
143        """Actionable message if ``src_mnt`` is NOT readable from inside WSL, else None.
144
145        Google Drive (``G:``/Drive File Stream), UNC and many network/virtual
146        filesystems are invisible to WSL's ``/mnt`` ? staging a video from such a path

```

```

145         fails with a cryptic ``cp: cannot stat`` and the whole run silently produces no
146         trajectory (i.e. "0 rallies", a wasted run that looks like a quality result).
147         A cheap ``test -r`` preflight turns that into one clear, fixable error BEFORE
148         any
149         frame extraction. Best-effort: a flaky/timed-out probe returns None so the real
150         ``cp`` still gets its chance (we never block a genuinely-readable source).
151         """
152         try:
153             res = self._wsl_bash(f"test -r {shlex.quote(src_mnt)} && echo READABLE")
154         except (subprocess.TimeoutExpired, OSError):
155             return None
156         if "READABLE" in (res.stdout or ""):
157             return None
158         return (
159             f"video source is not readable from WSL at '{src_mnt}'. Google Drive, UNC, "
160             f"and other network/virtual filesystems are invisible to WSL's /mnt ? copy "
161             f"the video to a local disk (e.g. C:\\ or the WSL filesystem) and retry."
162         )
163
164     class DetectorRunner(ABC):
165         """Common interface for shuttle/ball trajectory detectors (WSL or otherwise)."""
166
167         @abstractmethod
168         def run_predict(self, video_win_path: str, output_win_dir: str,
169                        video_id: Optional[str] = None) -> Optional[str]:
170             """Run inference and return path to trajectory CSV (or None on failure).
171
172             ``video_id`` (the file MD5) namespaces all staged/cached/output artifacts so two
173             videos that happen to share a filename cannot reuse each other's frames or CSV.
174             It is optional for back-compat; callers in the live pipeline always pass it.
175             """
176             pass
177
178         @abstractmethod
179         def healthcheck(self) -> Tuple[bool, str]:
180             """Quick check if the runner (env, weights, etc.) is usable.
181
182             Returns (ok, message). Never raises. Cheap to call before a long run.
183             """
184             pass
185
186         # De facto trajectory contract: the trajectory hybrids call these
187         # unconditionally after run_predict. Concrete runners reuse/re-export the
188         # model-agnostic TrackNetRunner staticmethods (same CSV schema + windowing);
189         # a runner with custom windowing supplies its own with the same shapes.
190         parse_trajectory_csv: Callable[..., Any]
191         trajectory_to_action_windows: Callable[..., Any]
192

```

8. user (2026-06-21T13:57:03.591Z)

C:\Users\avidu\Projects\khealsutra\docs\LOCAL_APPLIANCE_ARCHITECTURE.md

9. user (2026-06-21T13:57:03.737Z)

docs\DESKTOP_APP_PLAN.md

10. assistant (2026-06-21T13:57:07.758Z)

Let me note discrepancies so far on the base.py claim (lines 46-54 cited, but actual ONNX language is at 46-54 ? close). Now let me check the cited docs and the WASB-SBDT repo facts.

11. user (2026-06-21T13:57:08.755Z)

```

1      # Cross-Platform Desktop App ? "Plug in your card, get your highlights"
2
3      > **Status:** `DRAFT` ? `owner-gated` (scoping doc only; **nothing is built**) .
**Owner:** avidullu . **Created:** 2026-06-18 . **Last updated:** 2026-06-18
4      > **Lifecycle:** `DRAFT` ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
5      > **Tracking anchors:** $7 progress tracker is the source of truth; indexed in
`docs/README.md`; pointer in memory `current-state` (and `docs/NEXT_STEPS.md` once the owner
greenlights P0).
6      > **Relation to existing docs:** **peer-of**
[ `PLATFORM_ARCHITECTURE.md` ] (PLATFORM_ARCHITECTURE.md) (this is the **`LocalDesktop`
ComputeBackend** made into a product) . **realizes the L0 tier** of
[ `VIDEO_LOCALITY_MODEL.md` ] (VIDEO_LOCALITY_MODEL.md) . **consumes**
[ `OWNED_MODEL_IMPLEMENTATION_PLAN.md` ] (OWNED_MODEL_IMPLEMENTATION_PLAN.md) (the smaller/export-
clean model is one answer to the compute crux) . bound by
[ `COMMERCIALIZATION.md` ] (COMMERCIALIZATION.md) licensing guardrails.
7      > **Honesty note:** every claim is tagged `[verified]` (checked against code/measured),
`[researched]` (needs sign-off), or `[design]` (proposed). **The entire build is `[design]`** ?
this doc scopes; it does not ship. Not legal advice.
8
9      ---
10
11     ## 0. TL;DR
12     Package the mature local pipeline (WASB shuttle detector + improved windowing + ffmpeg
stitch, already runnable fully-offline via [ `tools/generate_highlights.py --skip-ai-
handoff` ] (./tools/generate_highlights.py)) as a **cross-platform desktop app** for **non-
technical recreational players**: plug in the camera's USB/SD card, click once, get back **only
the highlight reels** ? fully on-device, **no upload**, **no copy of the GBs of 4K originals**.
This is the product surface that fits the "local, cheap, efficient" thesis and the L0 ("nothing
leaves the machine") tier of `VIDEO_LOCALITY_MODEL.md`.
13
14     **The make-or-break unknown is compute on a typical enthusiast laptop** ? which often
has **no discrete GPU**. WASB runs at **~3.4x real-time on a laptop RTX 2070 Super** `[verified,
measured]`; on CPU it is far slower, and Apple Silicon has **no CUDA** at all. **P0 is a
feasibility spike that measures CPU and Apple-MPS throughput for a ~10-min clip before any app
is built.** If CPU-only is unacceptable on commodity hardware, the path needs a smaller/faster
owned model (ties `OWNED_MODEL_IMPLEMENTATION_PLAN`) or a tiered "fast-if-GPU, slow-but-works-
if-CPU" design. **Everything past P0 is gated on that result.**
15
16     ## 1. Problem & goal
17
18     **Why now / the user.** Target users are **recreational sports enthusiasts, not tech-
savvy**. Their match videos are **huge** (GBs of 4K). The cloud-upload product surface fails
them on every axis: slow on a typical Indian uplink (`VIDEO_LOCALITY_MODEL.md` $0: 10 GB ? ~1 h
at 20 Mbps, ~4.5 h at 5 Mbps), bandwidth-costly, privacy-fraught, and it bloats both our storage
and their drives. The thing the customer actually wants back is **kilobytes of timestamps ? a
highlight reel** (`VIDEO_LOCALITY_MODEL.md` $1: **"timestamps are the product; the video is dead
weight"**) .
19
20     **The concrete outcome ("good" looks like this):**
21     1. User plugs the GoPro/phone SD card or camera USB into their laptop.
22     2. The app detects the card, lists the match videos on it.
23     3. User clicks **Generate Highlights**.
24     4. The app makes an on-device proxy, runs WASB detection + windowing + stitches the reel
? **100% locally**.
25     5. The reel(s) land in a user-chosen folder. **The originals are never copied off the
card and never bloat any drive.**
26
27     No account, no upload, no cloud dependency in the default path. Gemini stays an
**optional** "better brain if you want it" toggle ? never required, never bundled as a hard
dependency, never a training source (`COMMERCIALIZATION` C4).
28
29     **Read to get current:** `VIDEO_LOCALITY_MODEL.md` (the L0/locality framing + OQ5
packaging seed), `PLATFORM_ARCHITECTURE.md` $3.2 (the ComputeBackend seam),
[ `backend/pipeline/detectors/base.py` ] (./backend/pipeline/detectors/base.py) (the detector ABC

```

that already anticipates a Linux-native runner), `tools/generate\_highlights.py` (the local entrypoint this app wraps).

```
30
31     ## 2. Decisions locked
32
33     | # | Decision | Source / date | Implication |
34     |---|-----|-----|-----|
35     | D1 | **Default path is 100% on-device, no upload, no Gemini.** The reel is produced
from the local pipeline (`--skip-ai-handoff`). | This doc . 2026-06-18; matches
`VIDEO_LOCALITY_MODEL` L0 | The app must never *require* a network call. Gemini is an opt-in
toggle only. |
36     | D2 | **Never copy the originals.** The app reads source video in place (off the card /
a chosen folder) and writes **only** reels + a tiny report to a user folder. | This doc; mirrors
`generate_highlights.py` non-destructive `_atomic_publish` discipline | No "import library" that
duplicates 10 GB files. Detection runs on a proxy; stitch reads the original in place. |
37     | D3 | **Drop the WSL2 dependency; run the detector natively per-OS.** | This doc;
enabled by the `DetectorRunner` ABC `[verified]` | Windows-native (no WSL), Linux+NVIDIA (CUDA
directly), macOS (MPS or CPU). The WSL adapter stays for the current dev box but is **not** the
shipped path. |
38     | D4 | **Commercial-clean licensing for everything bundled: MIT / Apache-2.0 / BSD
only** (code, weights). ffmpeg ships as an **LGPL build** (no GPL-only encoders like x264). |
`COMMERCIALIZATION` (ratified 2026-06-09); `VIDEO_LOCALITY_MODEL` Q05 ("ffmpeg LGPL build =
licensing-clean") | WASB-SBDT weights are MIT `[verified]` (see `tracknet-wasb-wsl2-decision`).
The stitch must use openh264 (BSD) or OS hardware encoders, **not** statically-linked x264. See
S3.4. |
39     | D5 | **Reuse the existing pipeline verbatim; add only the shell + a native runner +
device-I/O.** | `backend/pipeline/detectors/base.py` reuse contract `[verified]` | No re-
architecture of windowing/stitch. New code = native runner, USB/SD detection, app shell,
packaging. |
40     | D6 | **Feasibility is gated on P0.** No app shell, no packaging work starts until the
compute spike returns a verdict. | This doc; `decision-rigor-norm` (NUMBERS INFORM, FOUNDATIONS
DECIDE) | P0 is the only un-gated deliverable; P1?P3 are owner-greenlit *after* P0 numbers. |
41
42     ## 3. Foundation ? research / prior art `[design]` / `[researched]`
43
44     ### 3.1 The compute reality (the crux) ? what we know vs must measure
45     - **WASB on a discrete laptop GPU is acceptable** `[verified, measured]`: on the owner's
**RTX 2070 Super Max-Q (8 GB)**, WASB inference is the pipeline driver at **~3.2?3.4x real-
time** (?3.4 GPU-min per footage-min, ~0.055 s/frame), peak VRAM ~5.7?6.0 GB; proxy (NVENC)
~0.2x RT; 4K stitch (CPU libx264) ~16 s/rally; end-to-end ?5x clip length. Source: `current-
state.md` `COST_SCALING.md` checkpoint (6 Boxhill clips, 2026-06-18).
46     - **The typical enthusiast laptop has NO discrete GPU.** A 10-min 4K clip at 3.4x RT is
~34 GPU-min *with* a 2070S. On integrated graphics / CPU-only the wall-clock is unknown and
could be 10x+ worse ? possibly **hours per clip**, which would kill the one-click promise.
47     - **Apple Silicon has no CUDA.** WASB/TrackNet weights are PyTorch; PyTorch MPS (Metal)
backend *may* run them, but op-coverage and speed on the WASB model are **unverified**
`[design]`. CPU fallback on Apple is also unverified.
48     - **Owned-model tie-in.** `OWNED_MODEL_IMPLEMENTATION_PLAN.md` already states the on-
device answer: *train in Python ? **export across the boundary** (ONNX / Core ML / ExecuTorch /
GGUF) ? a thin device shell runs the artifact with zero Python*, with the standing constraint to
**train with export-clean ops**. A smaller, exported, quantized shuttle detector is a credible
answer if the CPU/MPS numbers on the current WASB weights are bad. **This makes P0 partly a
referendum on whether the desktop product needs the owned-model track.**
49
50     ### 3.2 The de-WSL native runner is already designed-for `[verified]`
51     The detector layer was deliberately built to make this port low-regret.
[`backend/pipeline/detectors/base.py`](../backend/pipeline/detectors/base.py) documents,
verbatim, an **"Optional Linux-native path (the main de-risk goal)"**: a
`LinuxNativeRunner`/`ONNXRunner` that **"Skip[s] all wsl / _stage / conda activate, load[s]
weights locally (torch or onnx)? write[s] the exact same CSV format so
`trajectory_to_action_windows` works verbatim."** The WSL plumbing lives in a separate
`WslCommandMixin`, explicitly so a native runner **must not inherit WSL plumbing.** The actual
inference math is already isolated in
[`wasb_infer.py`](../backend/pipeline/detectors/wasb_infer.py) (torch + the WASB-SBDT modules,
emitting `Frame,Visibility,X,Y`). **So the port is: lift `wasb_infer.py` out of WSL, load
```

```

weights with native torch/onnxruntime per-OS, keep the identical CSV contract.** The windowing
brain (`trajectory_to_action_windows`) and stitch are untouched.
52
53     ### 3.3 Compute backends to evaluate in P0
54     | Path | OS | Toolkit | Status |
55     |---|---|---|---|
56     | CUDA-native | Windows-native, Linux+NVIDIA | torch+CUDA (no WSL) | proven in WSL
`[verified]`; native = port |
57     | Apple MPS | macOS (Apple Silicon) | torch MPS / Core ML export | **unverified ? P0
must measure** |
58     | CPU | all (the no-GPU laptop) | torch-CPU / **onnxruntime** (often fastest on CPU) |
**unverified ? P0 must measure; the gating number** |
59     | Smaller/exported model | all | ONNX/Core ML/ExecuTorch, quantized | deferred to owned-
model track; P0 informs whether it's *needed* |
60
61     ### 3.4 App-shell prior art & licensing landscape `[researched]`
62     - **App shell options:** **Tauri** (Rust + system webview; ~3?10 MB binaries;
MIT/Apache) vs **Electron** (bundled Chromium; ~100?150 MB; MIT) vs **CLI + installer** (wrap
`generate_highlights.py`; smallest effort). All shell licenses are clean.
63     - **Backend packaging:** the Python/FastAPI backend bundles via **PyInstaller** or a
frozen venv as a local "sidecar." Packaging was already modernized ? **PEP 621 `pyproject.toml`
+ an `indexer` console entry point** shipped (`DEFERRED_HARDENING_PLAN` #7, DONE, PR #167)
`[verified]` ? so a CLI MVP has a real head start.
64     - **ffmpeg licensing (a genuine landmine):** ffmpeg is LGPL/GPL. The current stitch path
uses **libx264 (GPL)**. Redistributing that inside an installer would impose GPL on the bundle.
**Mitigation (D4):** ship an **LGPL ffmpeg build** using **openh264 (BSD)** or the OS hardware
encoder (NVENC / Apple **VideoToolbox** / Linux **VAAPI**) ? the desktop app should prefer
hardware H.264 anyway (faster, no CPU stitch tax).
65     - **Weights:** WASB-SBDT = MIT `[verified]` (`tracknet-wasb-wsl2-decision`),
redistributable.
66
67     ## 4. Design / architecture `[design]`
68
69     ### 4.1 The one-click flow (and what code each step reuses)
70     ```
71     [card inserted]
72     ? (NEW) per-OS device-insert watcher ? mount path
73     ?
74     [list videos on the card] ? glob video exts; reuse _VIDEO_EXTS from
generate_highlights.py
75     ?
76     [make on-device proxy] ? ffmpeg downscale (hardware encoder); reuse the proxy +
--detect-from split (#205)
77     ?
78     [detect: WASB native] ? (NEW) NativeWasbRunner (DetectorRunner subclass) ? NO WSL;
same CSV contract
79     ?
80     [windowing ? rallies] ? REUSE trajectory_to_action_windows + improved milestone
windowing (R1: cap=6 + R4)
81     ?
82     [stitch reel from ORIGINAL] ? REUSE VideoCompiler
(backend/pipeline/stitcher/compiler.py); read original in place
83     ?
84     [export reel + tiny report to a user folder] ? reuse _atomic_publish non-destructive
write; NEVER copy originals
85     ```
86     The detector runs on a **proxy**; the stitch reads the **original on the card** (the
`--detect-from` split, already shipped `[verified]` in `generate_highlights.py`). Rally
timestamps transfer 1:1 because the proxy shares the timeline (the collector frame-identical
proxy contract, `VIDEO_LOCALITY_MODEL` $1).
87
88     ### 4.2 Reuse map ? what's NEW vs REUSED
89     | Concern | Status | Code |
90     |---|---|---|
91     | Trajectory?rally windowing | **REUSE** `[verified]` |

```



```

`tracknet_runner.trajectory_to_action_windows` (model-agnostic staticmethod) |
92 | Stitch / reel compile | **REUSE** `[verified]` |
`backend/pipeline/stitcher/compiler.py::VideoCompiler` (pure ffmpeg) |
93 | Local orchestration (proxy?detect?stitch, non-destructive, resumable, `--skip-ai-
handoff`, `--detect-from`) | **REUSE** `[verified]` | `tools/generate_highlights.py` |
94 | Detector runner contract | **REUSE** `[verified]` |
`backend/pipeline/detectors/base.py::DetectorRunner` ABC |
95 | WASB inference math | **REUSE** (lift out of WSL)** `[verified]` |
`backend/pipeline/detectors/wasb_infer.py` |
96 | **Native (de-WSL) WASB runner, per-OS** | **NEW** `[design]` | `NativeWasbRunner` ?
sibling of `wasb_runner.py`, no `WslCommandMixin` |
97 | **USB/SD insert detection + video listing** | **NEW** `[design]` | per-OS: Win
WMI/`WM_DEVICECHANGE`, macOS DiskArbitration, Linux udev/udisks2 |
98 | **App shell / one-click UI** | **NEW** `[design]` | Tauri vs Electron vs CLI (§3.4; P2
decides) |
99 | **Packaging + code-signing per OS** | **NEW** `[design]` | PyInstaller/MSIX (Win),
.dmg+notarize (mac), AppImage/Flatpak (Linux) ? P3 |
100
101 ### 4.3 Where this sits in the platform architecture
102 This is the `local` / `LocalDesktop` ComputeBackend of
`PLATFORM_ARCHITECTURE.md` §3.2, realized as a shippable product instead of a server worker. The
control-plane/data-plane split degenerates cleanly to "all on one machine." Nothing here
forecloses the hosted SaaS path ? it's the privacy-max end of the same spectrum, and a genuine
differentiator (no competitor offers true on-device compute except SwingVision, which proves
it's commercially viable ? `PLATFORM_ARCHITECTURE` §intro).
103
104 ### 4.4 App-shell recommendation (to confirm at P2)
105 Lean: start with a CLI + installer MVP (wraps the existing
`indexer`/`generate_highlights.py` entry point ? fastest way to validate packaging, signing,
USB-detect, and the native runner end-to-end on all three OSes), then graduate to a Tauri
GUI for the non-technical user (tiny download fits the "huge-video laptops, slim app" frame;
the Python backend runs as a sidecar the Tauri front-end calls over localhost). Electron is the
fallback if webview inconsistencies bite. This mirrors the repo's "simplest thing first,
graduate on evidence" bias.
106
107 ## 5. Risk table `[design]`
108 | Risk | Severity | Mitigation / where decided |
109 |---|---|---|
110 | CPU/MPS too slow on a no-GPU laptop (kills the one-click promise) | make-or-
break | P0 spike measures it before anything is built (§6, §7-P0). Fallbacks: tiered
fast/slow UX, or the owned smaller model. |
111 | ffmpeg GPL contamination via libx264 in the installer | high (legal) | D4: LGPL build
+ openh264/hardware encoders only. |
112 | Code-signing friction (mac notarization mandatory; Win SmartScreen without EV cert) |
medium | P3 scopes certs/notarization per OS; budget the Apple Developer Program + a Windows
cert. |
113 | Privacy regressions (the on-device promise is the selling point) | medium | D1/D2:
no upload, no original-copy, no required network. Make the "nothing leaves your machine"
guarantee testable (assert no outbound calls in default mode). |
114 | Per-OS device-insert flakiness / removable-media edge cases (card pulled mid-run) |
medium | P1/P2 handle gracefully; never write to the card; fail safe. |
115 | Support burden for non-technical users (drivers, "it didn't detect my card") | medium
| favors a polished GUI + clear errors (reuse the actionable-error discipline already in the WSL
adapter). |
116
117 ## 6. Honest limits ? what this does NOT do
118 - A DRAFT/DONE status here proves nothing about runtime feasibility. The doc is a
scope; only the P0 numbers decide whether the product is buildable on commodity hardware.
119 - If CPU-only is too slow on a typical laptop, the simple "wrap the pipeline" path
does not exist ? it becomes "ship a smaller owned model" (a much larger effort, ties
`OWNED_MODEL_IMPLEMENTATION_PLAN`) or "GPU-tier only" (shrinks the addressable market). The doc
names this honestly rather than assuming the wrap is free.
120 - No accuracy claim changes. The desktop app inherits exactly the current rally-
detection quality (R1 milestone: precision is still the open gap per
`RALLY_DETECTION_GAP_CLOSURE.md`); packaging does not improve detection.

```

```

121 - **Gemini-quality reels need the optional online toggle.** The fully-local default uses
    `--skip-ai-handoff`; users wanting the Gemini brain accept that ?1280px chunks leave the machine
    (consent language, `VIDEO_LOCALITY_MODEL` OQ3).
122 - **macOS support is contingent on P0's MPS/CPU verdict** ? it is not assumed.
123
124 ## 7. Deliverables & progress tracker ? **source of truth**
125 Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. **One small PR per row**,
    branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.
    **Everything below is owner-gated; only this doc (the scoping PR) is in flight today.**
126
127 | ID | Deliverable | Depends on | Gated? | Status | PR |
128 |----|-----|-----|-----|-----|----|
129 | **D0** | **This scoping doc** (`DESKTOP_APP_PLAN.md` + README index) | ? | No | ? | |
    *(this PR)* |
130 | **P0** | **Compute-feasibility spike (make-or-break).** Measure WASB throughput for a
    ~10-min clip on: (a) CPU-only (torch-CPU & onnxruntime), (b) Apple MPS, (c) confirm CUDA-native
    (no WSL). Output: a wall-clock table + a go/no-go verdict + "do we need a smaller model?"
    recommendation. **No app code.** | D0 | **Owner greenlight** | ? | ? | |
131 | **P1** | **De-WSL native detector port** (`NativeWasbRunner`, per-OS): CUDA-native
    (Win/Linux), MPS/CPU fallback, identical `Frame,Visibility,X,Y` CSV; reuse windowing verbatim;
    healthcheck per backend. | P0 verdict | Yes | ? | ? | |
132 | **P2** | **Desktop app shell + one-click flow**: USB/SD detect ? list ? proxy ? native
    detect ? stitch ? export; shell choice (CLI MVP ? Tauri) confirmed; "no upload / no original-
    copy" enforced & testable. | P1 | Yes | ? | ? | |
133 | **P3** | **Packaging + code-signing installers** per OS (Win MSIX/Authenticode, mac
    .dmg + Developer-ID notarization, Linux AppImage/Flatpak); LGPL ffmpeg bundling; MIT/Apache/BSD
    audit of the bundle. | P2 | Yes | ? | ? | |
134
135 ## 8. Open questions ? owner / external
136 **Blocking gates (owner decides before the dependent phase):**
137 1. **(gates P1?P3) Greenlight P0** ? the spike is cheap and decides everything.
    Recommend: yes, run P0 next.
138 2. **(gates P1) If CPU is too slow, which fork?** tiered GPU-fast/CPU-slow UX · invest
    in a smaller exported owned model (`OWNED_MODEL_IMPLEMENTATION_PLAN`) · GPU-tier-only product
    (narrower market).
139 3. **(gates P2) App shell:** confirm CLI-MVP-then-Tauri, or go straight to a GUI /
    Electron.
140 4. **(gates P3) Distribution & signing budget:** Apple Developer Program (mandatory for
    notarization), a Windows code-signing cert (EV for clean SmartScreen) ? owner cost decision.
141
142 **Nice-to-knows:**
143 5. OS coverage priority for v1 (Windows-first? given the owner's box and the academy
    beachhead).
144 6. Auto-update strategy (Tauri updater / Sparkle / none for the pilot).
145 7. Does the desktop app reuse the optional Gemini toggle at all, or stay strictly local
    for v1?
146
147 ## 9. Definition of done
148 This subproject is DONE (? flip Status, archive) when:
149 - [ ] **P0** produced a measured throughput table (CPU / MPS / CUDA-native) and an
    owner-accepted go/no-go verdict (incl. whether a smaller model is required). *(P0 alone is a
    valid stopping point if the verdict is "not feasible without the owned model" ? then this doc
    hands off to `OWNED_MODEL_IMPLEMENTATION_PLAN`.)*
150 - [ ] **P1** the native (de-WSL) detector runs on ?1 GPU OS and ?1 CPU/MPS path,
    producing byte-compatible trajectories (regression vs the WSL runner on a shared clip).
151 - [ ] **P2** a non-technical user can insert a card and get a reel with no terminal, no
    upload, and no copied original ? verified end-to-end on ?1 OS.
152 - [ ] **P3** signed/notarized installers exist for the greenlit OSes; the bundle passes
    an MIT/Apache/BSD (+ LGPL-ffmpeg) license audit.
153 - [ ] $7 rows all ? and the headline feasibility verdict is recorded for posterity.
154
155 ## 10. References
156 **Internal code anchors `[verified]`:**
    [`tools/generate_highlights.py`](../tools/generate_highlights.py) (local orchestration; `--skip-
    ai-handoff`, `--wasb-stream`, `--detect-from`, `_WSL_DETECTORS`, `_atomic_publish`) ·

```

```
[`backend/pipeline/detectors/base.py`](../backend/pipeline/detectors/base.py) (`DetectorRunner`
ABC + the documented Linux-native de-risk path + `WslCommandMixin`) .
[`backend/pipeline/detectors/wasb_runner.py`](../backend/pipeline/detectors/wasb_runner.py) (the
WSL adapter to replace) .
[`backend/pipeline/detectors/wasb_infer.py`](../backend/pipeline/detectors/wasb_infer.py) (the
torch inference core to lift native) . [backend/pipeline/detectors/tracknet_runner.py`](../back
end/pipeline/detectors/tracknet_runner.py) (shared `trajectory_to_action_windows`) .
[`backend/pipeline/segmenters/wasb_hybrid.py`](../backend/pipeline/segmenters/wasb_hybrid.py) .
`backend/pipeline/stitcher/compiler.py::VideoCompiler` . `pyproject.toml` + `indexer` console
entry (#167).
157     **Internal docs:** [VIDEO_LOCALITY_MODEL.md](VIDEO_LOCALITY_MODEL.md) (L0 fully-local
tier; OQ5 packaging) . [PLATFORM_ARCHITECTURE.md](PLATFORM_ARCHITECTURE.md) §3.2
(ComputeBackend seam) .
[OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md) (export-clean on-
device model) . [COMMERCIALIZATION.md](COMMERCIALIZATION.md) (C4 paid-LLM no-train;
MIT/Apache/BSD-only) . [DEFERRED_HARDENING_PLAN.md](DEFERRED_HARDENING_PLAN.md) (#7 packaging,
DONE) . [TRACKNET_WSL_SETUP.md](TRACKNET_WSL_SETUP.md) (the WSL setup being superseded for
shipped builds).
158     **Measured source:** `current-state.md` 2026-06-18 `COST_SCALING.md` checkpoint (WASB
~3.4x RT on RTX 2070 Super Max-Q, peak VRAM ~5.7?6.0 GB).
159     **External (to validate at P0/P2/P3):** PyTorch MPS backend . onnxruntime
(CPU/CoreML/DirectML execution providers) . Tauri vs Electron sizing . ffmpeg LGPL build /
openh264 . Apple Developer-ID notarization . Windows Authenticode/MSIX.
160
161     ---
162
163     ### Changelog
164     - `2026-06-18` ? Created. Scoping doc for the cross-platform desktop app + the de-WSL
native-compute port; all `[design]`. §7 P0 (compute-feasibility spike) named the make-or-break
gate. Owner-gated; nothing built.
165
```

12. user (2026-06-21T13:57:10.524Z)

```
1:# Single-box appliance ? UI-powered rally indexer + reel generator (KhelSutra Studio)
16:Two honest hardware truths: (1) the current droplet is **CPU-only** ? real detection there is
**Gemini (cloud)** or **stub (CPU mock)**; **real local-GPU detection needs a GPU host** (native
WASB, in-flight). (2) The owner priority is to **stand up a GPU host ASAP** ? provider-agnostic
(DO GPU droplet / GCE GPU VM / Cloud Run+GPU) ? but that runs **in parallel** with the UI MVP,
which ships on Gemini today.
23:**Outcome ("good")** an operator signs in, picks a source, watches it index, curates
rallies, downloads a reel ? on the **live CPU droplet today** (Gemini/stub), and on a **GPU host
later** (native WASB) with **no UI fork**.
39:| D10 | Delivery = a **self-hosted webserver + UI the user spins up** (NOT a desktop app);
the user points at **local storage + GPU/CPU OR cloud locations** and the UI drives the tool |
Owner 2026-06-20 | **supersedes `DESKTOP_APP_PLAN.md`** ; one artifact (this appliance) runs on
the owner box, a droplet, or a GPU host ? same UI, compute/storage are settings |
43:- **In-process worker** one module-global `ThreadPoolExecutor(max_workers=1)` `[verified
job_worker.py:37]` ? strictly **one job at a time** ; the DB `jobs` table is the clock;
`fail_stale_jobs` sweeps queued/running on boot. Gemini chunk concurrency is a **separate** knob
`indexing.ai_max_workers` (**default 5** `[verified models.py:183]` ; set to 1 to serialize) ? it
is **not** already 1.
54:ONE box: **edge gate** (cloudflared+Access or Caddy basicauth) ? **reverse proxy / static
host** serving the `web/` SPA bundle and proxying `/api/` ? engine **FastAPI `127.0.0.1:8000`**
? **LocalStorage** ? **detector seam** (Gemini/stub today; native WASB on a GPU host). The
firewall blocks inbound `8000` + the test ports `9000/9001/9023`. The marketing site stays a
**separate** systemd unit on a **separate** subdomain.
131:| Surprise Gemini/GPU bill | **no cost ceiling/ledger** | per-video cap; `ai_max_workers`
(default 5) tunable; `max_workers=1` job serialization |
140:[Omitted long matching line]
150:- **L5 ? Owner hardware e2e (handed off).** Native WASB **byte-parity** on a real GPU box;
real-4K **golden-corpus** runs (Gemini + native); Modern-Standby/`cv2`-vs-`ffmpeg` behavior; a
multi-GB browser upload. Reported back into §9, stated plainly as **not verified** in the
offline suite.
173:**Owner priority (2026-06-20): T13 (GPU host) is elevated to run in parallel from the
```

```
start** ? owner provisions the host + token; the engine session lands native WASB; the UI MVP
does **not** block on it (ships on Gemini).
176:[Omitted long matching line]
180:- [ ] CUJ-1 works end to end on the droplet (Gemini **and** stub) behind the gate, via the
ingest menu.
194:[Omitted long matching line]
```

13. assistant (2026-06-21T13:57:16.359Z)

Now I have the local-codebase facts. Important discrepancy noted: claim 11 cites LOCAL_APPLIANCE_ARCHITECTURE.md "D10" as committing to "runs on Gemini today" appliance, AND DESKTOP_APP_PLAN.md "D3, D4 §3.1". But D10 actually says it **supersedes** DESKTOP_APP_PLAN.md, and the delivery is NOT a desktop app ? it's a self-hosted webserver. The claim's framing matches D10. Let me verify the wasb_infer.py to confirm the "torch + WASB modules" and check whether `requirements.txt` really pins `torch>=2.12.0` (it does ? line 9).

Now let me web-verify the external facts. Let me load WebSearch and WebFetch.

14. assistant (2026-06-21T13:57:21.285Z)

Let me verify the WASB-SBDT repo facts directly via the GitHub API (Dockerfile pins, LICENSE, src/ structure).

15. user (2026-06-21T13:57:23.816Z)

```
FROM nvcr.io/nvidia/tensorrt:21.06-py3
```

```
MAINTAINER shuheitarashima@ntt.com
LABEL OBJECT="Dockerfile for Widely Applicable Strong Baseline for Sports Ball Detection and
Tracking [BMVC2023]"
ARG http_proxy
ARG https_proxy
ENV http_proxy ${http_proxy}
ENV https_proxy ${https_proxy}
ENV DEBIAN_FRONTEND noninteractive
```

```
RUN apt-get update && apt-get install -y --no-install-recommends\
    apt-utils\
    apt-transport-https\
    ca-certificates\
    software-properties-common\
&& apt-get clean\
&& rm -rf /var/lib/apt/lists/*
RUN apt-get update && apt-get install -y \
    git\
    vim\
    cmake\
    libsm6\
    libxext6\
    libxrender-dev\
    libgl1-mesa-glx\
    locales\
    build-essential\
    zlib1g-dev\
    libssl-dev\
    libbz2-dev\
    libreadline-dev\
    libsqlite3-dev\
    libhdf5-dev\
    libncursesw5-dev\
    libgdbm-dev\
    liblzma-dev\
    libdb-dev\
    libffi-dev\
```

```
    uuid-dev\  
    tk-dev\  
&& apt-get clean\  
&& rm -rf /var/lib/apt/lists/*  
  
ENV LANG en_US.UTF-8  
ENV LANGUAGE en_US:en  
ENV LC_ALL en_US.UTF-8  
RUN locale-gen en_US.UTF-8  
  
RUN pip3 install numpy==1.22.4 torch==1.11.0+cu113 torchvision==0.12.0+cu113 --extra-index-url  
https://download.pytorch.org/whl/cu113  
RUN pip3 install hydra-core==1.2.0 tqdm==4.64.0 opencv-python==4.6.0.66 scikit-learn==1.1.1  
scikit-image==0.19.3 pandas==1.3.5 einops==0.4.1 timm==0.6.5 matplotlib==3.5.2  
  
COPY src /root/src  
WORKDIR /root/src  
#COPY .vimrc /root/
```

16. user (2026-06-21T13:57:25.081Z)

MIT License

Copyright (c) 2023 NTT Communications Corporation

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE

17. user (2026-06-21T13:57:25.953Z)

```
configs  
dataloaders  
datasets  
detectors  
losses  
main.py  
models  
optimizers  
runners  
setup_scripts  
trackers  
utils
```

18. assistant (2026-06-21T13:57:30.618Z)

Excellent ? claims 1, 2, and 3 confirmed directly from source. The Dockerfile confirms the exact pins (claim 1), MIT/NTT Communications Corporation 2023 (claim 2 ? note the claim says "NTT Communications" which matches "NTT Communications Corporation"), and src/ has configs/detectors/models/trackers (claim 3 ? plus dataloaders, exactly as claimed). Let me verify the models/ contents to confirm the "HRNet-style backbone via timm" claim, which is the

load-bearing part of claim 3.

19. user (2026-06-21T13:57:32.310Z)

```
__init__.py
ballseg.py
deepball.py
hrnet.py
monotrack.py
resnetv1b.py
resnet2d.py
segbase.py
unet2d.py
unet2d_parts.py
```

20. user (2026-06-21T13:57:41.021Z)

Exit code 134

```
=== __init__.py ===
=== ballseg.py ===
=== deepball.py ===
=== hrnet.py ===
=== monotrack.py ===
=== resnetv1b.py ===
=== resnet2d.py ===
=== segbase.py ===
=== unet2d.py ===
=== unet2d_parts.py ===
---grep timm/hrnet in models---
/usr/bin/bash: line 1: 1903 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1904 Broken pipe          | base64 -d 2> /dev/null
      1905 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1906 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1907 Broken pipe          | base64 -d 2> /dev/null
      1908 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1909 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1910 Broken pipe          | base64 -d 2> /dev/null
      1911 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1912 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1913 Broken pipe          | base64 -d 2> /dev/null
      1914 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1915 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1916 Broken pipe          | base64 -d 2> /dev/null
      1917 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1918 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1919 Broken pipe          | base64 -d 2> /dev/null
      1920 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1921 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1922 Broken pipe          | base64 -d 2> /dev/null
      1923 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1928 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1929 Broken pipe          | base64 -d 2> /dev/null
      1930 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
/usr/bin/bash: line 1: 1931 Done          gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1932 Broken pipe          | base64 -d 2> /dev/null
      1933 Aborted              (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
```

```
/usr/bin/bash: line 1: 1934 Done                gh api "repos/nttcom/WASB-
SBDT/contents/src/models/$f" --jq '.content' 2> /dev/null
      1935 Broken pipe                | base64 -d 2> /dev/null
      1936 Aborted                    (core dumped) | grep -il -e timm -e hrnet > /dev/null 2>&1
```

21. assistant (2026-06-21T13:57:44.215Z)

There's an `hrnet.py` in models/. Let me directly inspect hrnet.py and the default config to confirm the backbone and whether timm is actually the backbone source (the claim's specific wording).

22. user (2026-06-21T13:57:47.224Z)

```
=== ballseg.py head ===
```

23. assistant (2026-06-21T13:57:50.000Z)

The pipe to `head` is causing the base64 process to abort. Let me write to temp files instead.

24. user (2026-06-21T13:57:53.659Z)

```
=== hrnet.py imports ===
8:from __future__ import absolute_import
9:from __future__ import division
10:from __future__ import print_function
12:from __future__ import absolute_import
13:from __future__ import division
14:from __future__ import print_function
16:import os
17:import logging
19:import torch
20:import torch.nn as nn
=== models __init__ ===
from .unet2d import TrackNetV2
from .resunet2d import ChangsTrackNet
from .monotrack import MonoTrack
from .hrnet import HRNet
from .deepball import DeepBall
from .ballseg import BallSeg

__factory = {
    'tracknetv2': TrackNetV2,
    'monotrack': MonoTrack,
    'restracknetv2': ChangsTrackNet,
    'hrnet': HRNet,
    'deepball': DeepBall,
    'ballseg': BallSeg
}

def build_model(cfg):
    model_name = cfg['model']['name']
    if not model_name in __factory.keys():
        raise KeyError('invalid model: {}'.format(model_name))
    if model_name=='tracknetv2' or model_name=='resunet2d' or model_name=='monotrack':
        frames_in = cfg['model']['frames_in']
        frames_out = cfg['model']['frames_in']
        bilinear = cfg['model']['bilinear']
        halve_channel = cfg['model']['halve_channel']
        model = __factory[model_name]( frames_in*3, frames_out, bilinear=bilinear,
        halve_channel=halve_channel)
    elif model_name=='higher_hrnet' or model_name=='cls_hrnet' or model_name=='hrnet':
        model = __factory[model_name](cfg['model'])
    elif model_name=='restracknetv2':
        frames_in = cfg['model']['frames_in']
```

```

frames_out      = cfg['model']['frames_out']
halve_channel   = cfg['model']['halve_channel']
mode            = cfg['model']['mode']
neck_channels   = cfg['model']['neck_channels']
out_mid_channels = cfg['model']['out_mid_channels']
blocks          = cfg['model']['blocks']
channels        = cfg['model']['channels']
model = __factory[model_name]( frames_in*3,
                                frames_out,
                                mode=mode,
                                halve_channel=halve_channel,
                                neck_channels=neck_channels,
                                out_mid_channels=out_mid_channels,
                                blocks=blocks, channels=channels)

elif model_name=='deepball':
    frames_in      = cfg['model']['frames_in']
    frames_out     = cfg['model']['frames_out']
    class_out      = cfg['model']['class_out']
    block_channels = cfg['model']['block_channels']
    block_maxpools = cfg['model']['block_maxpools']
    first_conv_kernel_size = cfg['model']['first_conv_kernel_size']
    last_conv_kernel_size = cfg['model']['last_conv_kernel_size']
    first_conv_stride = cfg['model']['first_conv_stride']
    model = __factory[model_name]( frames_in*3, frames_out*class_out,
                                    block_channels=block_channels,
                                    block_maxpools=block_maxpools,
                                    first_conv_kernel_size=first_conv_kernel_size,
                                    first_conv_stride=first_conv_stride,
                                    last_conv_kernel_size=last_conv_kernel_size)

elif model_name=='ballseg':
    frames_in      = cfg['model']['frames_in']
    frames_out     = cfg['model']['frames_out']
    scale_factors = cfg['model']['scale_factors']
    backbone       = cfg['model']['backbone']
    model = __factory[model_name]( in_channels=frames_in*3, nclass=frames_out*1,
    backbone=backbone, scale_factors=scale_factors)

return model

```

25. assistant (2026-06-21T13:57:58.678Z)

Important nuance for claim 3. The WASB default model is **HRNet** (`hrnet.py`), but `hrnet.py` does NOT import `timm` ? it is a self-contained HRNet implementation (imports only `torch`). `timm` is used elsewhere (likely the `ballseg` backbone). So the claim's wording "HRNet-style backbone via `timm`" is slightly imprecise ? the default HRNet detector does not go through `timm`. But the broader claim (standard CNN, no exotic ops, `timm` is in deps, low-op-risk for `torch-2.x/ONNX`) holds. Let me confirm where `timm` is actually imported.

26. user (2026-06-21T13:58:01.890Z)

```

=== ballseg imports ===
=== default config (which model?) ===
dataloader
dataset
detector
eval.yaml
extract_frame.yaml
loss
model
optimizer
runner
tracker
transform

```


27. assistant (2026-06-21T13:58:05.680Z)

The pipe-to-head is killing base64 again on ballseg. Let me read it from the written file directly with Grep, and check the default model config.

28. user (2026-06-21T13:58:06.628Z)

```
<tool_use_error>Path does not exist: /tmp/ballseg.py. Note: your current working directory is C:\Users\avidu\Projects\badminton-highlight-indexer.</tool_use_error>
```

29. user (2026-06-21T13:58:09.916Z)

```
ballseg.yaml
deepball.yaml
deepball_large.yaml
monotrack.yaml
restracknetv2.yaml
tracknetv2.yaml
wasb.yaml
=== eval.yaml defaults ===
defaults:
  - _self_
  - runner: eval
  - dataset: badminton
  - model: wasb
  - dataloader: default
  - detector: tracknetv2
  - transform: default
  - tracker: online
hydra:
  run:
    dir: /root/outputs/${hydra.job.name}/${now:%Y-%m-%d_%H-%M-%S}
output_dir:
seed: 1234
```

30. assistant (2026-06-21T13:58:13.029Z)

The default model is `wasb` (model: wasb), detector `tracknetv2`, tracker `online`. Let me check what `wasb.yaml` model config points to (which backbone class) ? this is the actual shipped WASB model.

31. user (2026-06-21T13:58:16.764Z)

```
=== model/wasb.yaml ===
name: hrnet
frames_in: 3
frames_out: 3
inp_height: 288
inp_width: 512
out_height: 288
out_width: 512
rgb_diff: False
out_scales: [0]
MODEL:
  EXTRA:
    FINAL_CONV_KERNEL: 1
    PRETRAINED_LAYERS: ['*']
  STEM:
    INPLANES: 64
    STRIDES: [1,1]
  STAGE1:
    NUM_MODULES: 1
    NUM_BRANCHES: 1
    BLOCK: BOTTLENECK
```

```

    NUM_BLOCKS: [1]
    NUM_CHANNELS: [32]
    FUSE_METHOD: SUM
STAGE2:
    NUM_MODULES: 1
    NUM_BRANCHES: 2
    BLOCK: BASIC
    NUM_BLOCKS: [2,2]
    NUM_CHANNELS: [16,32]
    FUSE_METHOD: SUM
STAGE3:
    NUM_MODULES: 1
    NUM_BRANCHES: 3
    BLOCK: BASIC
    NUM_BLOCKS: [2,2,2]
    NUM_CHANNELS: [16,32,64]
    FUSE_METHOD: SUM
STAGE4:
    NUM_MODULES: 1
    NUM_BRANCHES: 4
    BLOCK: BASIC
    NUM_BLOCKS: [2,2,2,2]
    NUM_CHANNELS: [16,32,64,128]
    FUSE_METHOD: SUM
DECONV:
    NUM_DECONVS: 0
    KERNEL_SIZE: []
    NUM_BASIC_BLOCKS: 2
INIT_WEIGHTS: True

```

=== ballseg.py first 30 lines (via gh+python) ===

```

"""Image Cascade Network"""
import torch
import torch.nn as nn
import torch.nn.functional as F

from .segbase import SegBaseModel

"""
__all__ = ['ICNet', 'get_icnet', 'get_icnet_resnet50_citys',
           'get_icnet_resnet101_citys', 'get_icnet_resnet152_citys']
"""

class BallSeg(SegBaseModel):
    """Image Cascade Network"""
    def __init__(self, nclass = 19, backbone='resnet50', pretrained_base=False, in_channels=3,
scale_factors=[1, 1, 0.5]):
        super(BallSeg, self).__init__(nclass, in_channels=in_channels, backbone=backbone,
pretrained_base=pretrained_base)

        self._scale_factors = scale_factors

        self.conv_sub1 = nn.Sequential(
            _ConvBNReLU(in_channels, 32, 3, 2),
            _ConvBNReLU(32, 32, 3, 2),
            _ConvBNReLU(32, 64, 3, 2)
        )

```

32. assistant (2026-06-21T13:58:21.423Z)

This is a crucial finding for claim 3. The shipped WASB model (`model: wasb` ? `name: hrnet`) is a **custom HRNet built from a YAML spec of standard conv/bottleneck/basic blocks** ? defined directly in `hrnet.py` (imports only torch), NOT via timm. The `ballseg` model uses `segbase`/resnet backbone, also not timm in the file I read. Let me confirm where timm is

actually used in the repo (the claim says "timm==0.6.5 is the backbone source"), and check setup_scripts for timm usage to be fair to the claim.

33. user (2026-06-21T13:58:26.563Z)

```
=== python files in src ===
src/dataloaders/__init__.py
src/dataloaders/dataset_loader.py
src/dataloaders/heatmaps/__init__.py
src/dataloaders/heatmaps/heatmaps.py
src/dataloaders/img_transforms.py
src/dataloaders/samplers/__init__.py
src/dataloaders/samplers/samplers.py
src/dataloaders/seq_transforms.py
src/datasets/__init__.py
src/datasets/badminton.py
src/datasets/basketball.py
src/datasets/soccer.py
src/datasets/tennis.py
src/datasets/volleyball.py
src/detectors/__init__.py
src/detectors/deepball_detector.py
src/detectors/deepball_postprocessor.py
src/detectors/detector.py
src/detectors/postprocessor.py
src/losses/__init__.py
src/losses/bce.py
src/losses/combo_loss.py
src/losses/dice_loss.py
src/losses/focal_loss.py
src/losses/heatmap.py
src/losses/quality_focal_loss.py
src/losses/segmentation.py
src/losses/ssd_loss.py
src/losses/wbce.py
src/main.py
src/models/__init__.py
src/models/ballseg.py
src/models/deepball.py
src/models/hrnet.py
src/models/monotrack.py
src/models/resnetv1b.py
src/models/resnetv2d.py
src/models/segbase.py
src/models/unet2d.py
src/models/unet2d_parts.py
src/optimizers/__init__.py
src/optimizers/lr_schedulers.py
src/runners/__init__.py
src/runners/base.py
src/runners/eval.py
src/runners/extract_frame.py
src/runners/runner_utils.py
src/runners/train_and_test.py
src/trackers/__init__.py
src/trackers/intra_frame_peak.py
src/trackers/online.py
src/utils/__init__.py
src/utils/dataclasses.py
src/utils/evaluator.py
src/utils/file.py
src/utils/heatmap.py
src/utils/image.py
src/utils/refine_gt.py
src/utils/utils.py
```

```
src/utils/vis.py
```

```
=== searching code-search for timm in repo ===
```

```
1
Dockerfile
```

34. assistant (2026-06-21T13:58:33.272Z)

Decisive finding for claim 3: GitHub code-search reports `timm` appears in **only 1 file**: the Dockerfile (``total_count: 1`, item path = Dockerfile`). It is installed as a dependency but **not** actually imported anywhere in `src/`. The shipped WASB model is a self-contained HRNet (custom implementation using only torch). So the claim's statement "HRNet-style backbone via timm" and "timm==0.6.5 is the backbone source" is **incorrect on the mechanism** ? timm is an unused/transitive declared dep, not the backbone source. However, the claim's load-bearing conclusion (standard CNN, no exotic ops ? low-op-risk for torch-2.x port and ONNX export) is actually **strengthened** by this finding: the model is plain conv/bottleneck/basic-block HRNet. I'll mark this claim with a correction.

Now let me web-verify the external factual claims (4, 5, 6, 7, 8) in parallel.

35. user (2026-06-21T13:58:40.162Z)

Web search results for query: "conda-pack relocatable conda environment build OS must match target conda-unpack cannot relocate again license BSD"

Links: [{"title": "Conda-Pack ? conda-pack 0.8.1 documentation", "url": "https://conda.github.io/conda-pack/"}, {"title": "How to use conda-pack to relocate your conda environment? | by LITING CHEN | Medium", "url": "https://litingchen16.medium.com/how-to-use-conda-pack-to-relocate-your-conda-environment-622b68e077df"}, {"title": "Making packages relocatable ? conda-build 0.0.0.dev0+placeholder documentation", "url": "https://docs.conda.io/projects/conda-build/en/latest/resources/make-relocatable.html"}, {"title": "GitHub - conda/conda-pack: Package conda environments for redistribution · GitHub", "url": "https://github.com/conda/conda-pack"}, {"title": "Unpacked conda environment breaks if moved · Issue #106 · conda/conda-pack", "url": "https://github.com/conda/conda-pack/issues/106"}, {"title": "Conda-Pack Windows Instructions · GitHub", "url": "https://gist.github.com/pmbaumgartner/2626ce24adb7f4030c0075d2b35dda32"}, {"title": "Relocatable conda environments · Issue #1845 · conda/conda", "url": "https://github.com/conda/conda/issues/1845"}]

Based on the search results, here's what I found about conda-pack:

Key Information

conda-pack is a command line tool for creating relocatable conda environments, useful for deploying code in a consistent environment, potentially in locations where python or conda isn't already installed.

OS Compatibility

The OS where the environment was built must match the OS of the target, meaning that environments built on Windows can't be relocated to Linux.

Relocation Limitations

Once an environment is unpacked and conda-unpack has been executed, it cannot be relocated. Re-applying conda-pack is unlikely to work.

License

Conda-pack is offered under a New BSD license.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

36. user (2026-06-21T13:58:41.459Z)

Web search results for query: "micromamba static single binary BSD create environment conda-forge lock file no conda install"

Links: [{"title": "Micromamba Installation - Mamba's documentation!", "url": "https://mamba.readthedocs.io/en/latest/installation/micromamba-installation.html"}, {"title": "Advanced Usages ? micromamba-docker 2.6.0 documentation", "url": "https://micromamba-docker.readthedocs.io/en/latest/advanced_usage.html"}, {"title": "mamba - conda-forge | Anaconda.org", "url": "https://anaconda.org/conda-forge/mamba"}, {"title": "libmamba - conda-forge", "url": "https://prefix.dev/channels/conda-forge/packages/libmamba"}, {"title": "Mamba User Guide ? documentation", "url": "https://mamba.readthedocs.io/en/latest/user_guide/mamba.html"}, {"title": "New lock file format · Issue #1209 · mamba-org/mamba", "url": "https://github.com/mamba-org/mamba/issues/1209"}, {"title": "GitHub - mamba-org/micromamba-releases: Micromamba executables mirrored from conda-forge as Github releases · GitHub", "url": "https://github.com/mamba-org/micromamba-releases"}, {"title": "libmambapy - conda-forge | Anaconda.org", "url": "https://anaconda.org/conda-forge/libmambapy"}, {"title": "GitHub - mamba-org/mamba: The Fast Cross-Platform Package Manager · GitHub", "url": "https://github.com/mamba-org/mamba"}, {"title": "micromamba - conda-forge | Anaconda.org", "url": "https://anaconda.org/conda-forge/micromamba"}]

Based on the search results, here's what I found about micromamba, static binaries, BSD licensing, creating environments, conda-forge, and lock files:

Key Findings

****Static Single Binary:****

micromamba is a single-file executable that is statically linked and can be dropped anywhere on the operating system to get started with powerful package management and virtual environments.

****BSD License:****

The software is licensed under the BSD-3-Clause license.

****Creating Environments with conda-forge:****

You can create environments with micromamba using commands like ``micromamba create -n env_name xtensor -c conda-forge``

****Lock File Support:****

micromamba can be used to install lock files generated by conda-lock without having to install conda-lock. Simply invoke ``micromamba create -n my-env -f conda-lock.yml`` with an environment lockfile named `*-lock.yml` or `*-lock.yaml`.

****Benefits of Using Lock Files:****

When a lockfile is used to create an environment, the micromamba create command does not query the package channels or execute the solver, and using a lockfile has the added benefit of reducing the time to create a conda environment.

In summary, micromamba offers a lightweight, standalone solution for environment management with no conda installation required, making it particularly useful for containerized deployments and CI/CD pipelines.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

37. user (2026-06-21T13:58:41.677Z)

Web search results for query: "Anaconda terms of service defaults channel paid license 200 employees conda-forge free commercial"

Links: [{"title": "Update on Anaconda's Terms of Service for Academia and Research | Anaconda", "url": "https://www.anaconda.com/blog/update-on-anacondas-terms-of-service-for-academia-and-research"}, {"title": "Is Conda actually free? | pydevtools", "url": "https://pydevtools.com/handbook/explanation/is-conda-actually-free/"}, {"title": "Anaconda Legal | Anaconda.org", "url": "https://anaconda.org/conda-forge/micromamba"}]

```
Anaconda", "url": "https://www.anaconda.com/legal"}, {"title": "Anaconda Licensing Changes - by  
Charlotte Tolonen", "url": "https://blog.fulcrumgenomics.com/p/anaconda-licensing-  
changes"}, {"title": "Navigating Changes in Anaconda's Terms of Service for Miniconda's Commercial  
Users | by Bagus Tris Atmaja | Medium", "url": "https://medium.com/@btatmaja/navigating-changes-  
in-anacondas-terms-of-service-for-miniconda-s-commercial-  
users-5937efebb19b"}, {"title": "Navigating Anaconda Licensing Changes: What You Need to Know |  
DataCamp", "url": "https://www.datacamp.com/blog/navigating-anaconda-licensing"}, {"title": "How to  
Navigate Anaconda Licensing Challenges and Stay Compliant? | by Muthukrishnan |  
Medium", "url": "https://muthukrishnanpk84.medium.com/how-to-navigate-anaconda-licensing-  
challenges-and-stay-compliant-9a0b3d5aa3e1"}, {"title": "Anaconda packages are not  
?free?", "url": "https://sunpy.org/posts/2024/2024-08-09-anaconda/"}, {"title": "With Pro 3.1, can I  
avoid Anaconda license fees by using conda-forge?", "url": "https://community.esri.com/t5/arcgis-  
pro-questions/with-pro-3-1-can-i-avoid-anaconda-license-fees-by/td-p/1653595"}]]
```

Based on the search results, here's what I found regarding Anaconda's terms of service, the defaults channel, licensing requirements, the 200-employee threshold, conda-forge, and free commercial use:

Key Licensing Requirements

If you use the default channel (named "defaults" and points to the Anaconda Repository) and your company has 200 or more employees, you need to purchase a license.

The 200-Employee Threshold

The distribution of packages Anaconda builds and maintains remains free for individuals and small organizations, which is defined in the Terms of Service as having less than 200 employees or contractors. A paid license is required for large organizations, which is defined in the Terms of Service as having more than 200 or employees or contractors.

Conda-Forge and Free Commercial Use

If packages are hosted in community channels such as conda-forge, you can use all packages for free, provided you use Miniconda or any non-Anaconda installer such as Miniforge. Payment requirements in our Terms of Service do not apply to the user-uploaded packages and channels at anaconda.org, which includes the conda-forge channel.

Changes to Terms of Service

Anaconda introduced these commercial terms in April 2020, then clarified the 200-employee threshold in September 2020. A March 2024 update quietly removed those exemptions: nonprofits and government agencies above the threshold now require a paid license, and university research (for non-curriculum use) also requires licensing.

The key takeaway is that organizations can use conda and conda-forge for free regardless of size, but accessing Anaconda's "defaults" channel requires a paid license if the organization has 200 or more employees.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

38. assistant (2026-06-21T13:58:44.805Z)

Claims 4, 5, 6 confirmed. Now let me verify claims 6 (torch CUDA wheel sizes), 7 (PyTorch state_dict backward-compat), and 8 (ONNX Runtime execution providers + MIT license).

39. user (2026-06-21T13:58:51.383Z)

Web search results for query: "pytorch GPU CUDA wheel size 583MB bundled 4GB cudnn pyinstaller torch large"

Links: [{"title": "Why torch wheel is so huge (582MB)? · Issue #17621 ·"}]

pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/17621"}, {"title": "Streamline CUDA-Accelerated Python Install and Packaging Workflows with Wheel Variants | NVIDIA Technical Blog", "url": "https://developer.nvidia.com/blog/streamline-cuda-accelerated-python-install-and-packaging-workflows-with-wheel-variants/"}, {"title": "Missing Linux x86-64 Large wheels (cuda bundled) -- can not install alongside TensorFlow - PyTorch Forums", "url": "https://discuss.pytorch.org/t/missing-linux-x86-64-large-wheels-cuda-bundled-can-not-install-alongside-tensorflow/197651"}, {"title": "Help with Install Torch with Cuda - PyTorch Forums", "url": "https://discuss.pytorch.org/t/help-with-install-torch-with-cuda/211985"}, {"title": "torch · PyPI", "url": "https://pypi.org/project/torch/"}, {"title": "PyTorch 2.7 Release ? PyTorch", "url": "https://pytorch.org/blog/pytorch-2-7/"}, {"title": "Issues Installing PyTorch 2.7.0 and torchvision 0.17.0 with CUDA 12.9 on Windows 11 - Missing Wheel Files and CUDA Not Enabled - PyTorch Forums", "url": "https://discuss.pytorch.org/t/issues-installing-pytorch-2-7-0-and-torchvision-0-17-0-with-cuda-12-9-on-windows-11-missing-wheel-files-and-cuda-not-enabled/219527"}, {"title": "Issues Installing PyTorch 2.7.0 and torchvision 0.17.0 with CUDA 12.9 on Windows 11 - Missing Wheel Files and CUDA Not Enabled · Issue #383 · llyasviel/FramePack", "url": "https://github.com/llyasviel/FramePack/issues/383"}]

Based on the search results, here's what I found:

The PyTorch GPU wheel is approximately 583MB in size, which is quite large compared to other frameworks. This large size causes Docker images with PyTorch installed to typically be 2-4GB total, with PyTorch taking up at least 1/4 of that space.

Regarding bundled dependencies, the large wheel size is partially due to PyTorch being bundled with CUDA and cuDNN libraries. In smaller wheel variants, versions of CUDA libraries from PyPI are hardcoded, which can make it difficult to install alongside other frameworks like TensorFlow.

It appears that large bundled wheels were eventually replaced by smaller wheels in PyTorch versions after 2.2, though some users have requested that large wheels for PyTorch 2.2 and later be made available again.

Regarding PyInstaller specifically, the search results don't contain information about bundling PyTorch with PyInstaller in this context.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

40. user (2026-06-21T13:58:51.704Z)

Web search results for query: "pytorch state_dict backward compatibility load 1.11 into 2.x model recommended save load whole model not guaranteed"

Links: [{"title": "Implement pass-through `state\_dict` and `load\_state\_dict` for dynamo `OptimizedModule` · Issue #113422 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/113422"}, {"title": "torch.optim ? PyTorch 2.11 documentation", "url": "https://docs.pytorch.org/docs/2.11/optim.html"}, {"title": "Load Optimizer State Dict in PyTorch: A Comprehensive Guide ? codegenes.net", "url": "https://www.codegenes.net/blog/load-optimizer-state-dict-pytorch/"}, {"title": "module.load_state_dict doesn't load the state_dict modified for contiguous · Issue #98906 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/98906"}, {"title": "torch.compile can supports load_state_dict · Issue #129362 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/129362"}, {"title": "Mastering `model.load\_state\_dict` in PyTorch ? codegenes.net", "url": "https://www.codegenes.net/blog/modelupdate-state-dict-pytorch/"}, {"title": "BUG (potential crash) with `state\_dict()` implementation and overload · Issue #72778 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/72778"}, {"title": "Inconsistent model behavior when loading model using load_state_dict in different ways · Issue #23777 · pytorch/pytorch", "url": "https://github.com/pytorch/pytorch/issues/23777"}, {"title": "Changelog ? PyTorch Lightning 2.6.1 documentation", "url": "https://lightning.ai/docs/pytorch/stable/generated/CHANGELOG.html"}, {"title": "PyTorch 2.x", "url": "https://pytorch.org/get-started/pytorch-2-x/"}]

Based on the search results, I found some relevant information about PyTorch compatibility and state_dict handling:

PyTorch 2.0 is 100% backward compatible by definition, which suggests that basic compatibility between versions should be maintained. However, the search results don't contain specific documentation about loading state_dicts from PyTorch 1.11 into 2.x models or guarantees about loading whole models.

The search results do indicate general best practices: Be aware of PyTorch version compatibility, as some optimizer algorithms or their internal states may change between different versions, and it is recommended to use the same version of PyTorch when saving and loading the optimizer state dictionary.

However, the specific technical details you're asking about regarding backward compatibility guarantees for loading 1.11 state_dicts into 2.x models and whether saving/loading whole models is guaranteed versus recommended are not covered in these search results. For more detailed information on this topic, you may want to consult the official PyTorch migration guides or documentation specifically addressing version-to-version compatibility for state_dict loading.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

41. user (2026-06-21T13:58:53.436Z)

Web search results for query: "onnxruntime execution providers CPU CUDA DirectML CoreML MIT license run without pytorch"

Links: [{"title": "Execution Providers | onnxruntime", "url": "https://onnxruntime.ai/docs/execution-providers/"}, {"title": "Windows - DirectML | onnxruntime", "url": "https://onnxruntime.ai/docs/execution-providers/DirectML-ExecutionProvider.html"}, {"title": "Apple - CoreML | onnxruntime", "url": "https://onnxruntime.ai/docs/execution-providers/CoreML-ExecutionProvider.html"}, {"title": "Execution providers | ort", "url": "https://ort.pyke.io/perf/execution-providers"}, {"title": "Install ONNX Runtime | onnxruntime", "url": "https://onnxruntime.ai/docs/install/"}, {"title": "NVIDIA - CUDA | onnxruntime", "url": "https://faxu.github.io/onnxruntime/docs/execution-providers/CUDA-ExecutionProvider.html"}, {"title": "GitHub - microsoft/onnxruntime: ONNX Runtime: cross-platform, high performance ML inferencing and training accelerator", "url": "https://github.com/microsoft/onnxruntime"}, {"title": "Build with different EPs | onnxruntime", "url": "https://onnxruntime.ai/docs/build/eps.html"}, {"title": "CUDA Execution Provider - onnxruntime", "url": "https://oliviajain.github.io/onnxruntime/docs/execution-providers/CUDA-ExecutionProvider.html"}, {"title": "CoreML - onnxruntime", "url": "https://oliviajain.github.io/onnxruntime/docs/execution-providers/CoreML-ExecutionProvider.html"}]

Based on the search results, here's information about ONNX Runtime execution providers:

Execution Providers

ONNX Runtime supports multiple execution providers including NVIDIA CUDA, TensorRT, Intel OpenVINO and oneDNN, Windows DirectML, Qualcomm QNN, Android NNAPI, Apple CoreML, XNNPACK, AMD ROCm, MIGraphX, and Vitis AI.

Key Execution Providers

CPU: ONNX Runtime includes a CPU execution provider as a fallback option.

CUDA: The CUDA execution provider for ONNX Runtime is built and tested with CUDA 12.x and cuDNN 9.

DirectML: The DirectML Execution Provider is a component of ONNX Runtime that uses DirectML to accelerate inference of ONNX models. DirectML is only supported on Windows.

****CoreML****: The CoreML Execution Provider requires iOS devices with iOS 13 or higher, or Mac computers with macOS 10.15 or higher. It is recommended to use Apple devices equipped with Apple Neural Engine to achieve optimal performance.

Running Without PyTorch

For onnxruntime-gpu package, it is possible to work with PyTorch without the need for manual installations of CUDA or cuDNN. ONNX Runtime is an independent inference framework that doesn't require PyTorch to run?you can use it standalone with ONNX model files.

Regarding MIT License

I didn't find specific information about ONNX Runtime's MIT license in these search results. If you need clarification on the licensing, I'd recommend checking the official ONNX Runtime GitHub repository.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

42. assistant (2026-06-21T13:58:59.086Z)

Claims 6 and 8 substantially confirmed (torch wheel ~583MB ? 2-4GB bundled; ONNX Runtime supports CPU/CUDA/DirectML/CoreML and runs standalone without PyTorch). Two items need tightening:

- Claim 6 says "grows to ~4 GB bundled" and "two coexisting torch envs is roughly 10 GB" ? the search supports 2-4GB for a torch *Docker image*; the 4-5GB-per-env and 10GB-total are reasonable but on the high end. I'll mark confirmed-with-nuance.
- Claim 8's MIT license for onnxruntime wasn't confirmed by search. Let me verify the ONNX Runtime license directly and the PyTorch state_dict backward-compat doc.

43. user (2026-06-21T13:59:01.783Z)

MIT
MIT License

44. user (2026-06-21T13:59:03.119Z)

REDIRECT DETECTED: The URL redirects to a different host.

Original URL: https://pytorch.org/tutorials/beginner/saving_loading_models.html
Redirect URL: https://docs.pytorch.org/tutorials/beginner/saving_loading_models.html
Status: 301 Moved Permanently

To complete your request, I need to fetch content from the redirected URL. Please use WebFetch again with these parameters:

- url: "https://docs.pytorch.org/tutorials/beginner/saving_loading_models.html"
- prompt: "Does this page recommend saving/loading the state_dict over saving the whole pickled model? What does it say about backward compatibility of state_dict across PyTorch versions, and the risks of saving an entire model with pickle? Quote any relevant text."

45. assistant (2026-06-21T13:59:07.070Z)

onnxruntime MIT confirmed. Let me fetch the redirected PyTorch docs page to verify claim 7's state_dict guidance.

46. user (2026-06-21T13:59:13.892Z)

PyTorch's Recommendation on Model Saving

Yes, the page clearly recommends saving the state_dict over pickling the entire model.

Recommended Approach

For inference, the documentation states: "Saving the model's state_dict with the `torch.save()` function will give you the most flexibility for restoring the model later, which is why it is the recommended method for saving models."

Risks of Pickling Entire Models

The page explains several downsides of saving complete pickled models:

The serialized data becomes "bound to the specific classes and the exact directory structure used when the model is saved." This happens because pickle doesn't preserve the model class itself; instead, it "saves a path to the file containing the class, which is used during load time."

The consequence is significant: "your code can break in various ways when used in other projects or after refactors."

The documentation contrasts this with state_dict saving, which provides greater portability and robustness across different code versions and project structures.

47. assistant (2026-06-21T13:59:22.724Z)

Claim 7's core guidance is confirmed: state_dict is the recommended method; pickling the whole model breaks across refactors/versions. The claim's framing matches PyTorch's official guidance. The specific assertion "PyTorch officially guarantees state_dict backward-compat: a 1.11 state_dict loads into a 2.x model" is slightly stronger than what PyTorch formally guarantees ? a state_dict is just a dict of tensors keyed by parameter name, so it loads into an identically-architected module regardless of version (no formal cross-version "guarantee" per se, but practically robust for standard ops). I'll mark confirmed-with-nuance on the wording.

Now let me verify claim 9's specific line citations in native_wasb_runner.py and the byte-parity claim, and recheck claim 11's doc citation precisely (D10 vs DESKTOP_APP_PLAN). I already read both files. Let me verify the specific CSV header format and the wasb_infer torch usage to lock claim 9.

48. user (2026-06-21T13:59:24.852Z)

```
=== CSV header search ===
backend/pipeline/detectors/base.py:32:3. If it produces the same CSV schema
(Frame,Visibility,X,Y), reuse or re-export
backend/pipeline/detectors/native_wasb_runner.py:20:It emits the identical
``Frame,Visibility,X,Y`` CSV (same ``{stem}__{vid12}_wasb.csv``
backend/pipeline/detectors/rally_gate.py:10:we already produce (Frame,Visibility,X,Y). It
estimates the play axis (the image
backend/pipeline/detectors/stub_runner.py:3:A ``DetectorRunner`` that produces a deterministic
``Frame,Visibility,X,Y`` trajectory
backend/pipeline/detectors/stub_runner.py:11: * **replay** ? emit an existing
``Frame,Visibility,X,Y`` CSV verbatim (drive the
backend/pipeline/detectors/stub_runner.py:127: """CPU "mock GPU" for TRAINING: write a
deterministic Frame,Visibility,X,Y trajectory whose
backend/pipeline/detectors/tracknet_runner.py:9: 3. Read back the per-frame
``Frame,Visibility,X,Y`` CSV it writes.
backend/pipeline/detectors/tracknet_runner.py:238: """Parse a TrackNet predict CSV
(columns: Frame,Visibility,X,Y)."""
backend/pipeline/detectors/wasb_infer.py:8: Frame,Visibility,X,Y
backend/pipeline/detectors/wasb_runner.py:6: 3. runs it in the WSL `wasb` conda env -> a
Frame,Visibility,X,Y trajectory CSV,
```

49. assistant (2026-06-21T13:59:34.061Z)

All claim 9 elements verified: native_wasb_runner.py invokes `WASB\_PYTHON` by path (line 89), shells `wasb\_infer.py` (line 215), produces `Frame,Visibility,X,Y` CSV (line 20, wasb_infer.py:8); base.py:46-54 names "LinuxNativeRunner (or ONNXRunner)" that loads weights

"torch or onnx" (line 49) and writes "the exact same CSV format so trajectory_to_action_windows works verbatim" (line 50). Byte-parity 1194/1195 frames validated 2026-06-20 (lines 30-32). The line citations in the claim (52, 110-113, 210-237 + base.py:46-54) all match what I read.

Let me verify the precise location of the "583MB grows to 4GB" and reconfirm the LOCAL_APPLIANCE D10 wording vs the claim's framing ? the claim cites BOTH D10 (appliance, supersedes desktop) AND DESKTOP_APP_PLAN D3/D4 §3.1. There's tension: D10 supersedes DESKTOP_APP_PLAN. Let me read the exact D10 row context and the §3.1 / D3 / D4 I already have. I have enough. Let me also do a final check on the "google-genai is already a hard dep" recommendation premise ? confirmed in requirements.txt line 8 and pyproject line 35.

I now have everything needed. Let me compile the structured verdict.

50. user (2026-06-21T14:00:50.631Z)

Structured output provided successfully

51. assistant (2026-06-21T14:00:54.030Z)

I've completed the adversarial verification of all 12 claims. Summary: 11 confirmed, 1 (claim 3) confirmed on its load-bearing conclusion but refuted on a sub-mechanism, with the corrected claims and three actionable nuances captured in the structured output.