

claude-windows / dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3\subagents\workflows\wf\_f1431155-4a0\agent-a785ee2c14191391a.jsonl`
- SHA-256: `2a70c2745c8ca2aaf470e2d6f81b049729dab2b01c3e417520fe3278cb74b45f`
- Source modified: `2026-06-21T07:26:31+00:00`
- Imported at: `2026-07-05T16:48:22+00:00`
- project: `wf\_f1431155-4a0`
- session_id: `dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3`

Transcript

1. user (2026-06-21T07:23:57.135Z)

You are reviewing milestone M1 of the COMPUTE_DECOUPLED_SERVING project in the repo at C:/Users/avidu/Projects/badminton-highlight-indexer.
The spec is docs/COMPUTE_DECOUPLED_SERVING/README.md (read §2 D7-D15, §4.3 "Profile selection", and §7 the tracker row M1). M1 must be: PURE, ADDITIVE, with deployment.profile='' (the default) == today's behaviour, ZERO core change. A *deployment profile* is opt-in and applies a coherent preset bundle of EXISTING knobs (indexing.detector_impl/skip_ai_handoff, storage.backend) plus the new deployment.compute_target. Precedence: built-in defaults < profile preset < config.json < env.
D15: auto-detect SUGGESTS, never auto-applies (cloud spend never entered silently).

Files changed (read them all, plus the existing tests/test_config.py to understand the prior contract):

- backend/config/presets.py (NEW)
- backend/config/models.py (added DeploymentConfig + AppConfig.deployment field)
- backend/config/loader.py (added _deep_merge, _resolve_explicit_profile, preset application)
- backend/config/__init__.py (re-exports)
- tests/test_config_deployment.py (NEW)

Ground EVERY claim in the actual code with a file:line citation. Do NOT speculate. If you assert a bug, give the exact input that triggers it and the wrong-vs-right output. Rate confidence honestly; prefer FEWER, HIGH-CONFIDENCE findings. A clean bill ([] findings) is a valid and good result.

LENS: backward-compat / "zero core change". The hard requirement: profile='' yields a config byte-identical to before this PR. Verify by reading every caller surface: does adding the AppConfig 'deployment' section break anything that enumerates model_fields / dumps / snapshots config (grep the repo: backend + tests)? Does load_config now reading os.environ['DEPLOYMENT_PROFILE'] change behaviour when the var is unset? Does importing backend.config.presets at load time import torch or do any IO? Confirm tests/test_config.py's existing contracts (esp. test_config_shallow_section_replacement, test_non_object_config_degrades_to_defaults) still hold. Find any caller that constructs AppConfig with extra='forbid' or asserts an exact field set.

2. assistant (2026-06-21T07:23:59.953Z)

I'll review M1 of the COMPUTE_DECOUPLED_SERVING project. Let me start by reading the spec and the changed files.

3. user (2026-06-21T07:24:03.522Z)

Already up to date.

---BRANCH---

feat/serving-m1-deployment-profile

---STATUS---

M backend/config/__init__.py

M backend/config/loader.py

M backend/config/models.py

?? artifacts/

?? backend/config/presets.py

?? docs/KHELSUTRA_PRODUCT_OVERVIEW.md

?? tests/test_config_deployment.py

4. assistant (2026-06-21T07:24:05.482Z)

Let me read the spec and the changed files in parallel.

5. user (2026-06-21T07:24:06.984Z)

```
1      # Compute-Decoupled Rally Serving ? one pure core, profile-selected compute + storage
2
3      > **Status:** ? **M0 APPROVED** (owner, 2026-06-21) ? **IN PROGRESS (build M1+)** .
**Owner:** Avi . **Created:** 2026-06-21 . **Last updated:** 2026-06-21
4      > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to
`docs/archives/past_projects/` when DONE)
5      > **Tracking anchors:** §7 progress tracker is the source of truth; pointer in memory
`current-state` + `docs/NEXT_STEPS.md` once work starts.
6      > **Relation to existing docs:** the productionization successor to the archived
`DECOUPLED_COMPUTE` (M3/M4); realizes the `ComputeBackend`/`StorageBackend` registries from
[PLATFORM_ARCHITECTURE.md](../PLATFORM_ARCHITECTURE.md) §3; the **truly-local** profile is the
L0 tier of [VIDEO_LOCALITY_MODEL.md](../VIDEO_LOCALITY_MODEL.md) and the `LocalDesktop` backend
of [DESKTOP_APP_PLAN.md](../DESKTOP_APP_PLAN.md).
7      > **Naming:** this was the "Cloud Run + GPU serving subproject" (ADR-012). Renamed
**compute-decoupled serving** because **truly-local is a co-equal, first-class profile**, not an
afterthought.
8      > **Honesty note:** claims tagged `[verified]` (checked against code by the design
workflow `wf_56274ff0`), `[design]` (proposed), `[researched]`.
9
10     ---
11
12     > **? Northstar (D14):** a **mobile, app-like** flow ? point at a **Google Drive** video
? **fully cloud-native** run ? the stitched, **ready-to-share** reel lands **back in Drive, next
to the source**. Two design promises around it: a user can **switch local?cloud anytime** (D13,
a headline feature), and execution is **never a surprise** ? auto-detect *suggests*, the user
*confirms* (D15).
13
14     ## 0. TL;DR
15     The rally engine is a **pure function of `(input bytes + config) ? outputs`** in three
deterministic stages ? **DETECT ? WINDOW ? STITCH** ? that **no deployment profile may branch
on**. A **`ComputeBackend` (deployment profile)** decides *where* the core runs (a local process
vs a Cloud Run GPU job); a **`StorageBackend`** decides *where bytes come from* (local FS / USB
/ mounted Google Drive / bucket). **One startup config** picks both. We ship **truly-local**
(local GPU + in-process stitch, zero cloud) and **cloud-serving** (upload `<2GB` / gdrive /
bucket ? a Cloud Run GPU job runs detect **and** stitch, so stitch's compute is metered), plus
**cpu-mock** for CI. The single most important caveat: detection is **byte-identical only on the
same GPU** (cross-GPU is sub-pixel-different, per the 2026-06-20 deploy report) ? so the
**parity guarantee is enforced at the WINDOW + STITCH layer**, not detector-CSV bytes.
16
17     ## 1. Problem & goal
18     - **Why now:** DECOUPLED_COMPUTE proved compute *portability* (M0?M2 + M3.5 on a GCP L4)
and is archived; its deferred **M3 (serving endpoint + per-video billing) + M4 (scale-to-zero +
cost guard)** become this project. The owner needs **both** a hosted service *and* a truly-local
mode, with the **core (detect+stitch) unaffected** by which one runs.
19     - **The two user-facing outcomes** (owner, 2026-06-21):
```

20 - ***(a) Hosted service*** ? operate on small **local uploads (<2GB)** ***and*** load from **gdrive / a bucket**, spinning up **Cloud Run** for both rally segmentation AND stitching (stitching is compute-intensive ? run it where its cost is accounted for).

21 - ***(b) Truly-local*** ? on a local machine with videos on **local drives/USB + a local GPU**, run inference ***and*** stitching **completely locally**.

22 - **Good** looks like: the same **video ? rallies ? reel** core gives identical windows + stitched ranges whether it runs on a laptop or Cloud Run; switching is a **config/startup choice**, never a code branch in the core.

23

24 ## 2. Decisions locked

25

26 | # | Decision | Source | Implication |

27 |---|-----|-----|-----|

28 | D1 | The core is a **pure 3-stage function** (DETECT/WINDOW/STITCH); **no profile** branch inside it. | design **[verified]** | All profile difference lives in config + the two registries. |

29 | D2 | **Two registries:** **ComputeBackend** (where) selected by **deployment.profile** + **compute_target**; **StorageBackend** (what-bytes) by **input_backend/storage.backend**. | ADR-014 | Realizes **PLATFORM_ARCHITECTURE**'s **local/hosted/byo_worker**. |

30 | D3 | **Profiles shipped:** **truly-local**, **cloud-serving**, **cpu-mock** (CI); **byo_worker** reserved/deferred. | ADR-014 | Enum/seam reserved so BYO isn't precluded. |

31 | D4 | In **cloud-serving**, STITCH runs on Cloud Run (co-located with detect) so CPU/wall + egress are metered into the per-job cost ledger. | owner (a) | Stitch-on-laptop would hide real cost. |

32 | D5 | **Parity enforced at WINDOW+STITCH** (byte-exact for a fixed trajectory); **detect** is byte-exact only same-GPU. | deploy report 2026-06-20 | Cross-GPU asserted at the rally/window level, never CSV bytes. |

33 | D6 | **Default-OFF**, smallest-safe-first, one PR per milestone; any core-caller change (e.g. configurable output base) ships behind a flag + a Launch Report. | [[launch-report-convention]] | Zero-regression discipline. |

34 | D7 | **Cloud Run cold-start: SCALE-TO-ZERO** ? accept the ~1?3 min cold-start (it's an async job: upload?poll?reel, amortized over the multi-minute job). **No warm pool** (that ~\$500/mo idle would break per-video billing). | owner 2026-06-21 | M6/M7. Revisit a warm lane only on a UX complaint. |

35 | D8 | **Pricebook** = a versioned DB table; cost computed in **USD** (matches GCP billing = one source of truth), **displayed in INR** at a configurable FX rate (Bangalore market); re-versioned when GCP prices change. | owner 2026-06-21 | M8. |

36 | D9 | **Edge auth** = Cloudflare Access (reuse the site's existing CF). Lock the Cloud Run **ingress** to the CF proxy + **verify the Cf-Access JWT signature** (so a direct hit to the Cloud Run URL can't spoof the identity header). | owner 2026-06-21 | M9. One identity system. |

37 | D10 | **Free-tier "data-for-reels" trade:** free reels in exchange for **training rights** (uploaded footage + derived trajectories/labels); **paid tier** = no-train (privacy is the paid feature). **Carve-outs:** train ONLY on attested **ADULT** footage (minor footage ? reel yes, training pool **NO** ? DPDP/GDPR need verifiable parental consent); train on **DERIVED** data + **auto-delete raw**; **ToS counsel-reviewed**; live training-on-uploads **gated to M9** (public signup). Truly-local needs no DPA. | owner 2026-06-21 | M9 + D12; ties [[paper-coauthor-aim]] / PERSONALIZATION flywheel. |

38 | D11 | **Quality floor: Gemini handoff ON** for cloud-serving until the owned model clears the #172 ship-gate (e.g. ?0.70 multi-court) ? then flip via a **Launch Report**; the owned model runs in **shadow/eval** meanwhile. (Gemini = serving/inference ONLY, never a training source ? CLAUDE.md guardrail.) | owner 2026-06-21 | M7 + D12. |

39 | D12 | **Data-flywheel harness** is IN SCOPE (owner add, Q5): capture the (consented, adult) uploaded video + the Gemini reel/rally output ? **reliably store** in the per-video workspace/bucket ? **run shadow evals** (owned-vs-Gemini, dual-eval-set) ? **promote good ones** to the golden TRAINING set (human-reviewed labels) so the owned model climbs toward the D11 floor (then Gemini flips off). | owner 2026-06-21 | new **M11**; ties D10 (consent) + **data_pipeline/** + the eval/experiment harness. |

40 | D13 | **Profile** is user-switchable ANYTIME ? a headline **FEATURE**. The same user runs **local today, cloud tomorrow, back to local** ? it's just a config/startup choice; the pure core gives **equivalent** output either way. **Advertise it** ("your machine *or* our cloud ? your call, switch anytime"). | owner 2026-06-21 | Honest caveat: equivalent at the **rally/window** level, not bit-identical across GPUs (D5); a single job runs in one profile, switching is **between** jobs. |

41 | D14 | **NORTHSTAR** ? operate toward this: a **mobile, app-like** flow ? point at a

****Google Drive**** video ? ****fully cloud-native**** run ? the stitched, ****ready-to-share**** reel lands ****back in Drive, next to the source****. | owner 2026-06-21 | Implies a ****`gdrive-api`** (OAuth) **StorageBackend**** (read source + write reel back) ? distinct from the local ****mount**** backend; ****resolves \$8 #7****, new ****M12****. The mobile UI is the khelsutra track; the engine must support it (async jobs + status + Drive round-trip / signed URL). Net-new scope (OAuth + Drive write); not necessarily v1, but the design must not preclude it (the **`StorageBackend`** ABC accommodates it). |

42 | D15 | ****Auto-detect SUGGESTS, the user CONFIRMS ? execution is NEVER a surprise.**** A GPU owner may still choose cloud; ****cloud (= spend) is never entered silently****; the active profile is always explicit + surfaced. | owner 2026-06-21 | Amends §4.3 (auto-detect = a proposed default, not an auto-decision). |

43

44 **## 3. Foundation ? evolution / ADR lineage ? *trace the idea end-to-end***

45 This project is the latest step in a long decision chain; each ADR contributed a piece and a ****subtlety that carries forward****:

46

47 | ADR / doc | Contributed | Subtlety carried into this project |

48 |---|---|---|

49 | ****ADR-005**** | Permissive licensing (MIT/Apache-2.0/BSD only) | The engine + any served model + the ****container image**** (ffmpeg, fonts, weights) stay commercial-clean. |

50 | ****ADR-008**** | Owned-model topology; ****train==serve featurizer single-sourcing**** | The "core unaffected across profiles" is the same parity discipline, extended from train/serve to ****local/cloud****. |

51 | ****ADR-009**** | KhelSutra Guru; ****local-first delivery**** | The ****truly-local profile**** is the direct continuation ? privacy/offline self-host stays first-class. |

52 | ****ADR-010**** | DECOUPLED: D0?GCP pivot | The cloud-compute origin (GPU + co-located bucket). |

53 | ****ADR-011**** | DECOUPLED scope = M0?M2+M3.5; ****deferred M3 (serving) + M4 (billing/cost-guard)****; override the ****rent nothing / not a service**** posture | ****This project realizes M3/M4.**** Carries §06's owner-gated gates: ****DPA/consent for non-owner uploads, collector `md5` keying, WASB-C3****. |

54 | ****ADR-012**** | GPU-native; TPU deferred; ****Cloud Run+GPU scale-to-zero = serving model****; ****NVDEC**** named as the decode lever | The serving substrate + the 7-step seed scope this expands. |

55 | ****ADR-013**** | Drop D0; ****GCP sole compute stack****; \$200 D0 credits ? non-compute only | The provider for cloud-serving; D0 Spaces remains only a possible S3 ***storage*** target. |

56 | ****? ADR-014 (new)**** | ****Compute-decoupled serving:**** one pure core, profile-selected compute+storage; stitch-where-metered | This doc. ****Refines**** ADR-011/012, does ****not**** supersede them. |

57 | PLATFORM_ARCHITECTURE.md | The **`ComputeBackend`/`StorageBackend`** registries | The abstraction this realizes in code. |

58 | VIDEO_LOCALITY_MODEL.md / DESKTOP_APP_PLAN.md | L0 fully-local tier / **`LocalDesktop`** backend | The truly-local profile = these, productionized. |

59 | 07-COMPUTE-ABSTRACTION (archived) | **`DeviceContext`/`DeviceProfile`** (designed) | WASB hardcodes **`device='cuda'`** with ****no CPU fallback**** ? a portability gap M4 closes. |

60 | DEPLOY_REPORT_GCP_2026-06-20 (archived) | First real L4 run; ****cross-GPU sub-pixel parity finding**** | Why D5 (parity at window level, not CSV bytes). |

61

62 **## 4. Design / architecture**

63 See [**`architecture.svg`**](architecture.svg) for the picture. ****The core never changes; the box around it does.****

64

65 **### 4.1 The core contract (must stay byte/behaviour-identical) `[verified]`**

66 1. ****DETECT (GPU):**** **`DetectorRunner.run\_predict(video\_path, output\_dir, video\_id)`** ? CSV[Frame,Visibility,X,Y]` (**`backend/pipeline/detectors/base.py:164`**). Three interchangeable impls ? **`WasbRunner`** (WSL), **`NativeWasbRunner`** (Linux GPU), **`StubDetectorRunner`** (CPU mock) ? emit the ****identical CSV schema****. The trajectory CSV is the stable artifact boundary.

67 2. ****WINDOW (CPU pure fn):**** **`trajectory\_to\_action\_windows(points, fps, frame\_width, ?)`** ? **`windows`** (**`tracknet\_runner.py:281`**). A shared **`@staticmethod`**, zero GPU/IO/randomness ? same trajectory + config = byte-identical candidate windows. Used by every runner ***and*** the eval/regression stack verbatim.

68 3. ****STITCH (CPU/ffmpeg):**** **`VideoCompiler.compile\_highlights(raw\_video, segments, out\_name)`** ? reel` (**`stitcher/compiler.py:303`**). Pure FS+subprocess (merge ? ffmpeg ? watermark ? per-clip libx264 trim ? concat). Deterministic for a given input + ranges + watermark config.

69

```

70      **Non-goal (the trap to avoid):** *no* code branch inside detect/window/stitch keyed on
profile ? the same discipline the experiment harness already enforces (a disabled cue is byte-
identical to CONTROL).
71
72      ### 4.2 The deployment profiles
73      | Profile | Compute | Storage | Orchestration | When |
74      |---|---|---|---|---|
75      | **truly-local** | `compute_target=local_gpu`; `detector_impl=native` (Linux) / `auto`
(WSL); stitch **in-process** | `local` (or `gdrive` = mounted Drive) | `python -m backend.main`
or `worker infer` CLI; existing async job worker | Self-host / offline / privacy / dev-on-a-GPU-
box ? owner (b) |
76      | **cloud-serving** | `compute_target=cloud_run`; `detector_impl=native` (L4, cuda);
**detect AND stitch on the same Cloud Run job** | `gcs` (or gdrive mount / assembled upload);
per-video workspace; reels via signed URL | Cloud Run control plane enqueues ? Cloud Run GPU job
pulls/runs/pushes; `WAIT_AND_RESUME` reschedules the control-plane job; scale-to-zero | Hosted
SaaS ? owner (a) |
77      | **cpu-mock** | `compute_target=cpu_mock`; `detector_impl=stub` (synth/replay) |
`local`/in-memory fakes | `run_all_tests.py` / pytest | CI, regression, profile-parity, GPU-free
dev |
78      | *byo_worker* (deferred) | tenant GPU via outbound pull agent | tenant's own bucket
(signed URLs) | long-poll pull; same job table scoped by tenant | Privacy-sensitive enterprise ?
**DEFERRED**, enum/seam reserved |
79
80      ### 4.3 Profile selection (the "startup/setup config") `[design]`
81      ONE new top-level `deployment` section on `AppConfig` + preset application + env auto-
detect in `load_config`:
82      - `deployment.profile` ? {truly-local | cloud-serving | cpu-mock} (empty = today's
manual knobs, **fully backward-compatible**).
83      - Selecting a profile applies a **coherent preset bundle** of *existing* knobs that
today drift independently: **INPUT** (new `input_backend`/`input_root`, parallel to the output-
side `StorageConfig`), **COMPUTE** (new `compute_target` enum locking `detector_impl` +
`skip_ai_handoff` + workspace strategy), **STORAGE**
(`storage.backend`/`workspace_root`/`single_folder_per_video`, all already exist).
84      - **Precedence:** built-in defaults ? profile preset ? `config.json` ? env
(`RALLY_DETECTOR_IMPL`, `WASB_*`, `DEPLOYMENT_PROFILE`, ?) ? worker `--set k=v`. Every knob
stays individually overridable.
85      - **Env auto-detect SUGGESTS, the user CONFIRMS (D15 ? no surprise execution):** the env
gives a *proposed* default (`K_SERVICE` ? cloud-serving; `CI=true` ? cpu-mock; else
`torch.cuda.is_available()` ? truly-local), but the user **explicitly confirms or overrides** it
? a GPU owner may still pick cloud, and **cloud (= spend) is NEVER entered silently**. The
active profile is always surfaced (logged/shown). **Per-profile startup checks fail/warn fast**
(cloud-serving without GCS creds/GPU ? loud error before any job).
86      - The worker (`cmd_infer`/`cmd_train`) is **already profile-agnostic** (reads
`config.storage`, accepts `--config`/`--set`) ? no worker rewrite.
87
88      ### 4.4 Compute placement ? and *why stitch runs on Cloud Run* `[design]`
89      DETECT on GPU everywhere; WINDOW CPU-pure everywhere; STITCH is real CPU work. In
**cloud-serving, stitch is co-located in the Cloud Run job** so `gpu_active_seconds` (detect) +
`wall_seconds` (detect+stitch) + `bytes_out` (reel egress) land on **one metered invocation** ?
the per-job cost ledger. Running stitch on a laptop would leave that cost
**invisible/unattributed** ? exactly what the owner wants to avoid. (Stitch stays synchronous in
the GPU job for the common one-reel case; a queued CPU-only stitch fallback only if compiles get
bursty ? the job table + `claim_due_job` already supports both.)
90
91      ### 4.5 Reuse map ? most of the seams already exist `[verified]`
92      **Exists:** the 3-stage core; `_resolve_runner` (the single compute seam);
`StubDetectorRunner` (`replay_csv` = byte-exact $0 anchor); `StorageBackend` ABC + 4 backends +
`StorageRef.parse_uri`; the worker CLI (fetch/put, `--config`/`--set`); the async job control
plane (`jobs` table, `claim_due_job` CAS); `ExecutionPolicy WAIT_AND_RESUME`/`JobWaiting`
(carries to Cloud Run dispatch verbatim); the chunked-upload router (`<2GB`); `skip_ai_handoff`
(LocalNoAI); storage fakes; golden fixtures + experiment harness. **Partial:** per-video
workspace helper (default-OFF); `DeviceContext` (designed); edge-auth (v4 `user_id` columns
exist). **Net-new:** the `deployment` section; `input_backend`; configurable output base (kill
hardcoded `output`/); Cloud Run dispatch shim + container image; cost ledger + pricebook; edge-
auth header extraction; startup env detection.

```

```

93
94     ### 4.6 The data-flywheel harness (D12 / M11) `[design]`
95     Gemini-on (D11) ships good reels now; this harness is how we earn the right to flip
Gemini off. On a consented, adult cloud-serving job (the D10 trade), the pipeline
additionally:
96     1. Captures the source upload + the Gemini reel + the rally output
(intervals/report) into the per-video workspace (`workspaces/<video_id/>`), instead of the no-
consent auto-delete path.
97     2. Reliably stores them durably in the bucket (the consented-retention bucket; raw
video kept only for the training pool, not the no-consent path).
98     3. Runs shadow evals ? the owned model scores the same video alongside Gemini
(`backend/eval/experiment.compare` + the dual-eval-set, [[eval-dual-set-regression]]) ? tracks
the owned-vs-Gemini gap per job (the live read on "is the floor reachable yet?").
99     4. Promotes to golden ? captured videos that pass review become golden TRAINING
rows (human-reviewed labels via the annotation flow ? `data_pipeline/`), growing the corpus ?
the owned model climbs to the D11 floor ? flip Gemini off via a Launch Report.
100     This is the consent?capture?eval?goldenize flywheel (the cloud realization of the
PERSONALIZATION contribution model): every consented free reel makes the owned model better.
Reuses `data_pipeline/`, `backend/eval/`, and the M9 consent gate; adults-only, raw-
retention gated by the D10 consent.
101
102     ## 6. Honest limits ? what this does NOT do / prove
103     - A green CI suite proves plumbing + determinism (window+stitch) + profile-parity
for $0 ? it does NOT prove field quality, nor that real-GPU detection is good (that's owner-
side, real-video).
104     - Cross-GPU detection is NOT byte-identical ? parity is a window-level claim, not a
CSV-byte claim.
105     - The cloud `0.1.0-l4` weights are n=2 plumbing-grade ? cloud-serving ships with
Gemini handoff ON until the owned model clears a floor (open Q).
106     - This does not resolve WASB-C3 (sharing a trained model) or the DPA/consent
gate ? both remain owner/counsel-gated before public, non-owner serving.
107
108     ## 7. Deliverables & progress tracker ? source of truth
109     Legend: ? Todo · ? In progress · ? Done · ? Gated. One small PR per row, off
`origin/master`, no stacking. Default-OFF where it touches core callers.
110
111     | ID | Deliverable | Depends | Gated? | Status | PR |
112     |----|-----|-----|-----|-----|----|
113     | M0 | This design doc + project dir + `runlogs/` + ADR-014 | ? | owner sign-off | ? |
APPROVED 2026-06-21 |
114     | M1 | `deployment.profile` + `compute_target` config + preset bundles +
precedence/auto-detect; unit tests (no behaviour change, profile='' default) | M0 | safe | ? | ? |
|
115     | M2 | `input_backend`/`input_root`; wire main CLI/API input through
`StorageRef.parse_uri` (local=identity); storage-fake tests | M1 | safe | ? | ? |
116     | M3 | Configurable output/scratch base ? kill hardcoded `output/`
(`trajectory_hybrid:237,313`, `yolo_hybrid:407`, `gemini`); DEFAULT stays `output/`; golden
+ stub-E2E byte-identical | M1 | ? default-OFF + Launch Report on flip | ? | ? |
117     | M4 | `DeviceContext` + WASB CPU-fallback; unified healthcheck/device wiring | M1 |
safe (cuda default unchanged) | ? | ? |
118     | M5 | truly-local profile end-to-end (preset + startup checks); first committed
runlog (truly-local sample run) | M2,M3,M4 | owner real-GPU | ? | ? |
119     | M6 | Cloud Run GPU dispatch shim (control plane ? job via storage ? re-claim) +
containerized worker image (ffmpeg+CUDA+WASB+fonts); mocked-dispatch tests | M5 | ? owner
(GCP spend) | ? | ? |
120     | M7 | cloud-serving profile e2e on L4: upload`<2GB` + gdrive/bucket ? detect+stitch
on Cloud Run ? reel via signed URL; DEPLOY_REPORT (posterity, sample runs + contact sheet) |
M6 | ? owner (spend) | ? | ? |
121     | M8 | Per-job cost ledger (v_ migration: `owner_id, gpu_active_seconds,
wall_seconds, bytes_out, cost_usd, cost_breakdown_json`) + pricebook + aggregation +
`/api/.../cost` | M7 | ? owner (billing) | ? | ? |
122     | M9 | Edge auth (Cf-Access extraction) + per-tenant scoping on `user_id`;
DPA/consent gate wiring | M7 | ? owner (privacy/counsel) | ? | ? |
123     | M10 | `byo_worker` / zero-infra-BYO ? design-only, DEFERRED | M8,M9 | ? explicit
owner approval | ? | ? |

```

```

124 | **M11** | **Data-flywheel harness (D12):** on a consented adult job, **capture** the
upload + Gemini reel/rally output ? **reliably store** under the per-video workspace ? **shadow-
eval** owned-vs-Gemini (dual-eval-set / `experiment.compare`) ? **promote** good ones to the
golden TRAINING set (human-reviewed labels via the annotation flow). Closes the loop so the
owned model climbs to the D11 floor. Reuses `data_pipeline/` + `backend/eval/` + the consent
gate (M9). | M7,M9 | ? owner (consent/privacy) | ? | ? |
125 | **M12** | **gdrive-api` (OAuth) StorageBackend (D14 Northstar):** cloud reads the
source from + writes the stitched reel **back to the user's Google Drive next to the source**
(per-user OAuth consent) ? the mobile/cloud-native Drive round-trip. Distinct from the local
mount backend; slots into the `StorageBackend` ABC. | M7,M9 | ? owner (OAuth/consent) | ? | ? |
126
127 **Testing strategy is a first-class deliverable** ?
[ `TESTING_STRATEGY.md` ](TESTING_STRATEGY.md). **Run logs / sample runs for posterity** ?
[ `runlogs/` ](runlogs/).
128
129 ## 8. Open questions ? owner *(blocking-gates vs nice-to-knows)*
130 **? The 5 gating questions are LOCKED (owner, 2026-06-21) ? see §2 D7?D12:**
131 1. ? **Cold-start ? D7** (scale-to-zero, accept cold-start).
132 3. ? **Pricebook ? D8** (versioned DB table; USD internal / INR display).
133 4. ? **Edge auth ? D9** (Cloudflare Access; verify Cf-Access JWT, lock ingress to
proxy).
134 5. ? **DPA/consent ? D10** (free-tier data-for-reels trade; adults-only; derived-data;
counsel-gated to M9).
135 6. ? **Quality floor ? D11** + the **data-flywheel harness D12 / M11** (Gemini-ON until
the owned model clears the floor; capture?store?eval?goldenize meanwhile).
136
137 **Still open (non-gating, smaller):**
138 2. **Stitch placement:** always co-located in the detect job (simplest, one metered
unit) vs split to a cheap CPU Cloud Run service when bursty? *Default co-located unless metrics
say otherwise.*
139 7. ? **RESOLVED ? D14:** cloud needs a **`gdrive-api` (OAuth) StorageBackend** (read the
source from + write the reel back to the user's Drive) ? the Northstar's Drive round-trip ?
distinct from the local **mount** backend (which stays for truly-local). New milestone **M12**.
140 8. **NVDEC/NVENC** in the Cloud Run image (cut stitch wall-time/cost) vs libx264
(determinism/parity)?
141 9. **Input cap:** is `<2GB` a hard client reject, or accept to the 4GB upload cap and
route larger files elsewhere?
142 10. **Directory name:** `COMPUTE_DECOUPLED_SERVING/` (chosen) vs keep
`CLOUD_RUN_GPU_SERVING/` ? confirm.
143
144 ## 9. Definition of done
145 - [ ] M1?M4 land (config + input + output-base + device layer), each green on golden
cross-OS + experiment no-regression.
146 - [ ] **Profile-parity test green** ? truly-local vs cloud-mock produce byte-identical
windows + segment ranges.
147 - [ ] **truly-local** (M5) and **cloud-serving** (M7) each have a committed **runlog /
DEPLOY_REPORT** with a real sample run.
148 - [ ] Cost ledger (M8) makes GPU + **stitch** + egress cost visible per job; edge auth
(M9) gates multi-tenant.
149 - [ ] **Data-flywheel harness (M11)** captures?stores?shadow-evals?goldenizes consented
adult jobs; the owned model is on a measured path to the D11 floor.
150 - [ ] ? flip Status DONE, archive to `docs/archives/past_projects/`.
151
152 ## 10. References
153 **ADR lineage:** ADR-005/008/009/010/011/012/013 ? **ADR-014**
(`docs/archives/decisions/DECISIONS.md`). **Strategy:** `PLATFORM_ARCHITECTURE.md`,
`HOSTING_PLAN.md`, `VIDEO_LOCALITY_MODEL.md`, `DESKTOP_APP_PLAN.md`,
`LOCAL_APPLIANCE_ARCHITECTURE.md` (khelsutra). **Archived predecessor:**
`docs/archives/past_projects/DECOUPLED_COMPUTE/` (esp. `05-COST-ATTRIBUTION.md`, `07-COMPUTE-
ABSTRACTION.md`, `06-RISKS-AND-OPEN-QUESTIONS.md`, `DEPLOY_REPORT_GCP_2026-06-20.md`). **Code
anchors `[verified]`:** `backend/pipeline/detectors/base.py:164`, `tracknet_runner.py:281`,
`stitcher/compiler.py:303`, `segmenters/trajectory_hybrid.py:65,237`, `backend/storage/`,
`backend/worker/__main__.py:88`, `backend/api/{uploads,job_worker}.py`,
`backend/infrastructure/database.py`. **Design provenance:** workflow `wf_56274ff0` (4 code-
mappers ? synthesis; 58 findings).

```

```

154
155 ---
156 ### Changelog
157 - 2026-06-21 ? created from the design workflow output; M0 in progress.
158

```

6. user (2026-06-21T07:24:07.074Z)

```

1      """Deployment-profile presets + environment auto-detection (COMPUTE_DECOUPLED_SERVING
M1).
2
3      A *deployment profile* is a single switch that applies a coherent BUNDLE of existing
4      config knobs, so the truly-local / cloud-serving / cpu-mock modes don't drift apart one
5      knob at a time (the design in ``docs/COMPUTE_DECOUPLED_SERVING/README.md`` §4.3).
6
7      Precedence (lowest ? highest), realised by the loader::
8
9          built-in model defaults < profile preset < config.json < env < worker --set
10
11     This module is **pure data + pure functions** ? no IO, and (critically) **no torch
import at
12     module load**. The loader applies a preset ONLY when a profile is EXPLICITLY selected
13     (``config.json`` ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); with no
14     profile selected the loader is byte-identical to the pre-M1 behaviour (``profile=""`` ==
15     today). Auto-detect only ever **suggests** a profile ? it never auto-applies one (D15:
cloud
16     spend is never entered silently).
17     """
18
19     from __future__ import annotations
20
21     import copy
22     import os
23     from typing import Any, Mapping, Optional
24
25     # Canonical enum values, defined ONCE here. ``models.DeploymentConfig``'s ``Literal``
mirrors
26     # these and ``tests/test_config_deployment.py`` pins parity so the two can never drift.
27     PROFILES: tuple[str, ...] = ("truly-local", "cloud-serving", "cpu-mock")
28     COMPUTE_TARGETS: tuple[str, ...] = ("local_gpu", "cloud_run", "cpu_mock")
29
30     # Each profile maps 1:1 to its canonical compute_target. Used to apply the preset and to
warn
31     # when config.json explicitly overrides compute_target into a state inconsistent with
the profile.
32     COMPUTE_TARGET_FOR_PROFILE: dict[str, str] = {
33         "truly-local": "local_gpu",
34         "cloud-serving": "cloud_run",
35         "cpu-mock": "cpu_mock",
36     }
37
38     # A preset is a partial, nested override dict over the model defaults. It touches ONLY
knobs
39     # that EXIST today ? M1 is pure/additive: the new
``deployment.{profile,compute_target}`` plus
40     # ``indexing.{detector_impl,skip_ai_handoff}`` and ``storage.backend``. Later milestones
extend
41     # these bundles ? ``input_backend`` (M2) and the configurable output/workspace base (M3,
42     # Launch-Report-gated) ? deliberately NOT set here so M1 never pre-empts a gated flip.
43     PROFILE_PRESETS: dict[str, dict[str, Any]] = {
44         "truly-local": {
45             "deployment": {"profile": "truly-local", "compute_target": "local_gpu"},
46             # detector_impl stays "auto" (the runner picks WSL vs native by env on the GPU
box).
47         "Zero-cloud / privacy-first ? run fully local with no Gemini handoff.

```

```

48         "indexing": {"skip_ai_handoff": True},
49         "storage": {"backend": "local"},
50     },
51     "cloud-serving": {
52         "deployment": {"profile": "cloud-serving", "compute_target": "cloud_run"},
53         # Linux L4 (cuda); Gemini handoff ON until the owned model clears the #172 floor
(D11).
54         "indexing": {"detector_impl": "native", "skip_ai_handoff": False},
55         "storage": {"backend": "gcs"},
56     },
57     "cpu-mock": {
58         "deployment": {"profile": "cpu-mock", "compute_target": "cpu_mock"},
59         # Deterministic CPU mock ? no GPU/WSL, no provider/API. The CI / regression
profile.
60         "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
61         "storage": {"backend": "local"},
62     },
63 }
64
65
66 def is_known_profile(profile: str) -> bool:
67     """True for a non-empty, recognised profile name."""
68     return bool(profile) and profile in PROFILE_PRESETS
69
70
71 def preset_for(profile: str) -> dict[str, Any]:
72     """Return a deep COPY of the preset bundle for ``profile`` (so callers can't mutate
the
73     module-level table), or ``{}`` for ``""`` / an unknown profile."""
74     return copy.deepcopy(PROFILE_PRESETS.get(profile, {}))
75
76
77 def _truthy(val: Optional[str]) -> bool:
78     """Common env-var truthiness (``1/true/yes/on``, case/space-insensitive)."""
79     if val is None:
80         return False
81     return val.strip().lower() in {"1", "true", "yes", "on"}
82
83
84 def _cuda_available() -> bool:
85     """Best-effort CUDA probe. torch is heavy + optional, so it is imported lazily here
and
86     ONLY when a caller opts into ``probe_cuda``; any import/runtime failure means 'no
GPU'."""
87     try:
88         # Intentionally lazy + local: keeps the config-load path torch-free (torch is
heavy
89         # and optional). Only reached when a caller opts into probe_cuda.
90         import torch
91
92         return bool(torch.cuda.is_available())
93     except Exception:
94         return False
95
96
97 def suggest_profile(
98     env: Optional[Mapping[str, str]] = None, *, probe_cuda: bool = False
99 ) -> str:
100     """SUGGEST (never auto-apply ? D15) a deployment profile from the environment.
101
102     Signals, in priority order:
103
104     - ``K_SERVICE`` set (Cloud Run)           ? ``"cloud-serving"``
105     - ``CI`` truthy                            ? ``"cpu-mock"``
106     - ``probe_cuda`` and a CUDA GPU present ? ``"truly-local"``

```

```

107         - otherwise                ? ``"`` (no strong suggestion)
108
109     ``probe_cuda`` defaults to ``False`` so the common config-load path stays cheap and
never
110     imports torch; a CLI ``--detect`` or cloud-serving startup that wants the GPU probe
passes
111     ``probe_cuda=True``. The result is advisory only: the caller decides whether to act
on it.
112     """
113     e = os.environ if env is None else env
114     if e.get("K_SERVICE"):
115         return "cloud-serving"
116     if _truthy(e.get("CI")):
117         return "cpu-mock"
118     if probe_cuda and _cuda_available():
119         return "truly-local"
120     return ""
121

```

7. user (2026-06-21T07:24:08.037Z)

```

1     """Pydantic models for the application configuration.
2
3     This module defines the canonical, validated, typed configuration for the
4     Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6     The models support graceful migration from old config.json files via extra="allow".
7     """
8
9     from __future__ import annotations
10
11     from typing import Any, List, Literal, Optional
12
13     from pydantic import BaseModel, Field, field_validator, model_validator
14
15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0
29         max_scene_cuts_per_minute: float = 0.5
30         relevance: ValidationRelevanceConfig =
Field(default_factory=ValidationRelevanceConfig)
31
32
33     class TrackNetConfig(BaseModel):
34         wsl_distro: str = "Ubuntu"
35         repo_dir: str = "~/models/TrackNetV4"
36         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
37         conda_env: str = "TrackNetV4"
38         weights_path: str = ""
39         queue_length: int = 5
40         stage_in_wsl: bool = True
41         wsl_stage_dir: str = "~/clips"
42         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
43         velocity_threshold: float = 0.02

```

```

44     # Cache hygiene: after a SUCCESSFUL run the staged video copy (and, for WASB, the
45     # decoded frame cache) are deleted automatically ? they are pure intermediates,
46     # regenerable from the source, and the dominant disk consumer (~GBs/video). Set
47     # true only for debugging. On FAILURE they are always kept so a re-run can resume.
48     keep_frames: bool = False
49     # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB.
50     # 0 = disabled.
51     min_free_gb: float = 0.0
52
53
54     class WasbIndexConfig(BaseModel):
55         """Typed config for the live WASB shuttle substrate (indexing.wasb).
56
57         Mirrors backend/pipeline/detectors/wasb_runner.py::WasbConfig.from_indexing_cfg so
58         these knobs actually take effect ? previously this whole block was silently dropped
59         because nested config sections defaulted to extra='ignore'.
60         """
61         wsl_distro: str = "Ubuntu"
62         repo_dir: str = "~/models/WASB-SBDT"
63         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
64         conda_env: str = "wasb"
65         weights_path: str = "~/models/WASB-
SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
66         sport: str = "badminton"
67         wsl_stage_dir: str = "~/clips"
68         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
69         velocity_threshold: float = 0.02
70         # --- Linux-native runner only (detector_impl='native'); ignored by the WSL runner.
71         ---
72         # Env vars override these (WASB_DEVICE / WASB_PYTHON) for a config-free GPU box.
73         device: str = "cuda" # torch device for the native runner: cuda | mps | cpu
74         python_bin: str = "python" # the `wasb` conda env python (invoked by path, no
`conda activate`)
75         # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
76         # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates,
77         # regenerable
78         # from the source. Set true only for debugging. On FAILURE they are kept so a re-run
79         # resumes.
80         keep_frames: bool = False
81         # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB. 0 =
82         # disabled.
83         min_free_gb: float = 0.0
84         # FAST PATH ? decode the video on the fly and feed frames straight to the detector,
85         # skipping the slow per-frame PNG extraction that dominates a long clip (~46k PNGs
86         # for a
87         # 12-min 1080p60 video, which overran the 30-min timeout). Output is bit-identical
88         # to the
89         # PNG path (PNG is lossless) ? VERIFIED on a real GPU 2026-06-20 (native stream ==
90         # WSL disk,
91         # 1194/1195 frames byte-identical; native run-to-run byte-identical /
92         # deterministic).
93         # Tri-state:
94         # None = per-runner default ? the WSL runner keeps DISK extraction
95         # (unchanged Windows
96         # behaviour); the Linux-native runner (GPU box) defaults to STREAMING
97         # (the perf win).
98         # True/False = force on/off for BOTH runners (e.g. set False for a container cv2
99         # can't seek).
100         stream_video: Optional[bool] = None
101
102     class FusionConfig(BaseModel):
103         """Config for the multi-signal `fusion_hybrid` segmenter (MULTI_SIGNAL_FUSION_PLAN
104         P2).
105

```

```

95     Every default reproduces control behaviour: the fusion segmenter persists the
96     shuttle net-crossing count (Gate A signal) and ? only when ``cue_flags['person']``
is
97     true ? the person-cue features, but it does NOT gate the served handoff unless a
98     threshold is raised (``min_crossings``/``min_players`` default 0 = OFF). The court
mask
99     is optional: ``court_polygon=None`` ? Gate B counts every detection (player-count-
only);
100     supplied ? off-court persons (spectators / adjacent courts) are excluded.
101     ""
102     # Which trajectory substrate seeds the windows (shuttle stays primary).
103     substrate: Literal["wasb", "tracknet"] = "wasb"
104     # Windowing velocity threshold for the fusion family (the engine reads
105     # indexing.<family>.velocity_threshold). Default matches the substrates' 0.02; set
106     # equal to indexing.wasb.velocity_threshold to A/B fusion against a tuned wasb run.
107     velocity_threshold: float = 0.02
108     # Per-cue enable flags, e.g. {"person": true}. Person cue is OFF unless explicitly
109     # enabled ? so the default path never imports torch / touches the GPU.
110     cue_flags: dict[str, bool] = Field(default_factory=dict)
111     # Served Gate A: drop windows whose shuttle net-crossings < this from the AI
handoff.
112     # 0 = OFF (no gating; persisted candidates are always the full superset for offline
re-scoring).
113     min_crossings: int = Field(default=0, ge=0)
114     # Served Gate B: when the person cue is on, drop windows with median in-court
players
115     # < this. 0 = OFF.
116     min_players: int = Field(default=0, ge=0)
117     # Optional court mask as normalised 0-1 [[x, y], ...] polygon. None = no mask.
118     court_polygon: Optional[list[list[float]]] = None
119     # Net line for the person two-sided-motion split (person features only ? the shuttle
120     # net-crossing gate keeps rally_gate's camera-agnostic auto-estimate).
121     net_axis: Literal["x", "y"] = "y"
122     net_position: float = Field(default=0.5, ge=0.0, le=1.0)
123
124
125     class IndexingConfig(BaseModel):
126         default_segmenter: str = "yolo_hybrid"
127         use_gpu: bool = True
128         # Detector backend for the trajectory hybrids (wasb/tracknet). "auto" = the real
129         # WSL runner (default, production); "stub" = a deterministic CPU mock (no GPU/WSL)
130         # so the whole pipeline is regression-testable on a CPU box / CI; "native" = the
131         # Linux-native WASB runner (no WSL/conda-activate; a GPU box ? M3.5). The
132         # RALLY_DETECTOR_IMPL env var overrides this. See
pipeline/detectors/{stub,native_wasb}_runner.py.
133         detector_impl: str = "auto"
134         motion_threshold: float = 0.08
135         min_segment_duration: float = 3.0
136         dead_time_merge_gap: float = 2.0
137         chunk_duration: float = 90.0
138         chunk_overlap: float = 10.0
139         detector_model: str = "fasterrcnn_resnet50_fpn"
140         detector_score_threshold: float = 0.5
141         yolo_velocity_threshold: float = 0.15
142         yolo_frame_skip: int = 3
143         yolo_tracker: str = "bytetrack.yaml"
144         yolo_prefetch_workers: int = 2
145         min_action_window_duration: float = 2.0
146         # BOUNDED WINDOWING (over-merge fix W1, RALLY_QUALITY_RESEARCH.md §6). 0 = OFF; > 0
=a window
147         # longer than this many seconds is recursively split at its largest internal
inactivity gap (?
148         # ``window_min_split_gap`` s ? the most likely rally boundary), so one window can't
swallow
149         # several rallies + the dead time between them. Never merges, only cuts (recall

```

```

can't drop).
150      # *** R1 MILESTONE DEFAULT-ON (cap=6): the smallest safe local-windowing flip.
Measured n=13
151      # golden, R2 tolF1: baseline?milestone (cap=6 + R4 below) ?+0.222, sign 12/0,
p=0.0005; cap=6
152      # beats cap=12 ?+0.110, sign 12/0; net over-seg guard PASS (mean merge_rate
0.651?0.161). ***
153      max_window_duration: float = 6.0
154      window_min_split_gap: float = 0.5 # min internal gap (s) that qualifies as a split
point
155      # HARD-CAP fallback for W1: when True, an over-long window with NO internal
inactivity gap
156      # (a continuously-active trajectory ? e.g. the real-footage mahadevpura-1 174s blob)
is
157      # chopped at its midpoint until ? max_window_duration, making the cap a true upper
bound.
158      # Only cuts (recall can't drop). OFF by default until measured/promoted.
159      window_hard_cap: bool = False
160      # R4 rally-END inactivity trim (s, 0 = OFF). After a rally the shuttle keeps moving
slowly
161      # (picked up / walked) above velocity_thresh, so the window over-extends its END
(the measured
162      # |?end| gap vs golden). Trim the post-rally tail back to the last frame whose
trailing
163      # window stays dense with FAST motion (velocity > inactivity_rally_thresh). Only
cuts.
164      # *** R1 MILESTONE DEFAULT-ON (1.5/0.20/0.30): on top of cap=6, R4 adds a small but
significant
165      # end-precision gain ? ?tolF1 +0.040, sign 9/1/3, p=0.022, |?end| 4.96?3.31s (n=13).
The W1 cap
166      # is the dominant lever; R4 is the marginal increment. See QUALITY_ITERATIONS.md
"R1". ***
167      window_inactivity_end: float = 1.5
168      inactivity_rally_thresh: float = 0.20
169      inactivity_min_density: float = 0.30
170      # R5 blob-fallback (default OFF). LAST-RESORT splitter for the continuously-active
blob: after
171      # the W1 gap-split AND the R4 inactivity trim have each had their chance, a group
STILL longer
172      # than window_blob_factor × max_window_duration (no internal gap, no rally-end
signal ? e.g.
173      # mahadevpura-1's ~65s residual ? 11× a 6s cap) is midpoint-chopped to ? the cap as
a FORCED cut
174      # (logged + flagged on the window) so the degenerate single-window outcome is
avoided and the
175      # clip is surfaced as needing rally-STATE cues, not a bigger cap. Differs from
window_hard_cap
176      # (which chops BEFORE R4 and over-segments normal clips): blob-fallback runs AFTER
R4 and only
177      # force-cuts the genuine blob (the factor gate). Only cuts (recall can't drop).
178      window_blob_fallback: bool = False
179      # Magnitude gate that keeps R5 surgical: only a group longer than this ×
max_window_duration is a
180      # blob (a normal long rally ~1-2× the cap is left alone; the mahadevpura-1 residual
is ~11×).
181      # Default 4.0 is do-no-harm on the golden corpus (rescues only the continuously-
active
182      # clips, every other clip a no-op within noise) ? see docs/QUALITY_ITERATIONS.md
"R5".
183      window_blob_factor: float = 4.0
184      log_yolo_candidates: bool = True
185      store_yolo_candidates: bool = True
186      ai_handoff_padding: float = 2.0
187      ai_max_workers: int = 5
188      # Width (px) the gemini segmenter re-encodes chunks to before upload. Stream-copied

```

```

189     # chunks of a high-bitrate (4K) source blow past the provider upload cap
190     # (backend/providers/gemini.py::MAX_UPLOAD_MB) and are refused ? observed live as
191     # "0 segments / success". Re-encoding also makes chunk boundaries frame-accurate
192     # (stream copy cuts at the nearest keyframe). 0 = legacy stream-copy at source res.
193     # Matches gemini_refine's --clip-scale-width default.
194     ai_chunk_scale_width: int = 1280
195     # FULLY-LOCAL mode (default OFF): when true, the trajectory-hybrid segmenters
196     # (wasb_hybrid / tracknet_hybrid / fusion_hybrid) skip the Phase-3 AI handoff
entirely and
197     # emit their local candidate windows AS the rallies ? no provider, no API key,
nothing
198     # uploaded. Lets the local detector run standalone (e.g. to benchmark it against the
Gemini
199     # path, or to operate with zero cloud dependency). The pure `gemini` segmenter
ignores this.
200     skip_ai_handoff: bool = False
201     # Optional debug knob: trim every video to the first N seconds before processing.
202     debug_duration: Optional[float] = None
203
204     tracknet: TrackNetConfig = Field(default_factory=TrackNetConfig)
205     wasb: WasbIndexConfig = Field(default_factory=WasbIndexConfig)
206     fusion: FusionConfig = Field(default_factory=FusionConfig)
207
208     @field_validator("default_segmenter")
209     @classmethod
210     def segmenter_must_be_registered(cls, v: str) -> str:
211         # Registry check happens at runtime in main/segmenter selection.
212         # Here we just allow any string for forward compatibility.
213         return v
214
215     @model_validator(mode="after")
216     def _chunk_step_must_be_positive(self) -> "IndexingConfig":
217         # The chunked segmenters advance by (chunk_duration - chunk_overlap); a non-
positive
218         # step makes `while t < duration: t += step` loop forever, growing memory before
any
219         # GPU work. Reject the bad config at load time instead.
220         if self.chunk_overlap >= self.chunk_duration:
221             raise ValueError(
222                 f"indexing.chunk_overlap ({self.chunk_overlap}) must be < chunk_duration
"
223                 f"({self.chunk_duration}); otherwise chunking never advances."
224             )
225         return self
226
227
228     class OutputConfig(BaseModel):
229         default_highlights_name: str = "highlights.mp4"
230         db_name: str = "sports_indexer.db"
231         # Directory where the DB, highlight reels, and processed clips are written.
232         # The CLI's --output-dir overrides this at startup (see backend/main.py:main).
233         output_dir: str = "output"
234
235
236     class WatermarkConfig(BaseModel):
237         """Branding overlay burned into each highlight clip during the per-clip re-encode
238         (backend/pipeline/stitcher/compiler.py). **On by default** ? set `enabled: false`
239         (under `stitching.watermark` in config.json) to turn it off. The text is fully
240         configurable. The concat stage stays stream-copy (-c copy)."""
241         enabled: bool = True
242         text: str = "Generated by Khelsutra.guru"
243         logo_path: str = "" # optional transparent PNG overlay; "" = no logo
244         font_path: str = "" # optional .ttf; "" = use the fontconfig `font` family
below
245         font: str = "Sans" # fontconfig family used when font_path is empty

```

```

246     font_size_ratio: float = Field(default=0.035, gt=0.0, le=0.5) # text height as a
fraction of frame height
247     opacity: float = Field(default=0.85, ge=0.0, le=1.0)
248     box: bool = True # faint backing box behind the text for legibility
249     margin: int = Field(default=24, ge=0) # px inset from the chosen corner
250     position: Literal["bottom-right", "bottom-left", "top-right", "top-left"] = "bottom-
right"
251
252
253     class StitchingConfig(BaseModel):
254         begin_padding: float = 0.5
255         end_padding: float = 0.5
256         use_cache: bool = False
257         min_shot_count: int = 1
258         # Long-rally reel filter (seconds; 0 = OFF). Drop detected rally pairs whose
duration
259         # (end ? start) is below this from the highlight reel, BEFORE compiling. A cleaner
260         # reel-selection knob than `min_shot_count`: shot_count is unlabeled / "Unknown" on
the
261         # fully-local no-AI path, whereas duration is derived from the rally's own
start/end. The
262         # local detector over-fires and its false positives are mostly SHORT (motion blips,
263         # background-court fragments), so this keeps the reel to the eye-catching long
rallies.
264         # OFF by default ? the detector output is unchanged, only which rallies reach the
reel.
265         min_rally_duration: float = Field(default=0.0, ge=0.0)
266         watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
267
268
269     class LoggingConfig(BaseModel):
270         """Logging knobs for the run entrypoints (see backend/utils/logging_setup.py)."""
271         # Third-party HTTP/RPC client loggers (httpx, httpcore, urllib3, google-genai,
272         # google.auth, grpc) emit one INFO line per request ? dozens per Gemini run, which
273         # buries our own [rally]/[compile] lines in run.log during manual QA. Default raises
274         # them to WARNING; set true to restore full chatter for request-flow debugging.
275         verbose_http: bool = False
276
277
278     class SecurityConfig(BaseModel):
279         # When set, /api/process video_path must resolve under this directory (hosted/tenant
280         # mode). Default None preserves the single-user local 'any local file' workflow.
281         allowed_video_root: Optional[str] = None
282
283         # --- Chunked upload (B2, PR-2) ---
284         # Landing zone for browser uploads (assembled from parts). Per-user subdirectories
285         # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
286         # contain upload_dir or /api/process will reject the uploaded paths.
287         upload_dir: str = "output/uploads"
288         # Whole-file cap (default 4 GB) and per-part cap. Parts stay under ~100 MB because
289         # free-tier edge proxies cap request bodies there (HOSTING_PLAN §2).
290         max_upload_mb: int = 4096
291         max_part_mb: int = 95
292         allowed_upload_extensions: List[str] = ["mp4", "mov"]
293
294
295     class StorageConfig(BaseModel):
296         """Pluggable storage backend (docs/DECOUPLED_COMPUTE/). Default ``local`` = today's
297         behaviour (filesystem paths, byte-for-byte). ``s3`` covers AWS + Cloudflare R2 + D0
Spaces
298         + MinIO (via ``endpoint_url``); ``gcs`` = Google Cloud Storage; ``gdrive`` = a
locally MOUNTED
299         Google Drive (``gdrive.mount_root``). Per-backend settings are loose dicts passed to
the constructor ?
300         **secrets are never stored here**, only endpoints / regions / file paths / source-

```

```

selectors
301         (boto3/google auth chains supply credentials). See ``backend/storage/``.
302         backend: Literal["local", "s3", "gcs", "gdrive"] = "local"
303         # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
304         output_root: str = ""
305         # --- Per-video workspace (docs/PER_VIDEO_WORKSPACE.md) ---
306         # One folder per video keyed by the MD5 video_id:
<root>/[<workspace_prefix>]/<video_id>/...
307         # so a local deploy and a cloud (bucket) deploy share ONE relative layout (only the
root
308         # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5).
When OFF
309         # the legacy flat layout is used unchanged.
310         single_folder_per_video: bool = False
311         # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
312         workspace_root: str = ""
313         # Cloud namespace prefix under the root so per-video folders sit tidily beside
314         # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
directly.
315         workspace_prefix: str = "workspaces"
316         local: dict[str, Any] = Field(default_factory=dict)
317         s3: dict[str, Any] = Field(default_factory=dict)
318         gcs: dict[str, Any] = Field(default_factory=dict)
319         gdrive: dict[str, Any] = Field(default_factory=dict)
320
321
322     class DeploymentConfig(BaseModel):
323         """Deployment profile ? one switch that selects a coherent bundle of compute +
storage
324         knobs (COMPUTE_DECOUPLED_SERVING M1, ``docs/COMPUTE_DECOUPLED_SERVING/README.md``
$4.3).
325
326         **Pure / additive:** the empty default profile (``""``) applies NO preset, so the
pipeline
327         behaves exactly as before this section existed. A profile is opt-in (``config.json``
328         ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); auto-detect only ever
329         *suggests* one (D15 ? cloud spend is never entered silently). Nothing in the core
reads
330         these fields yet (M1): selecting a profile works purely by pre-setting the EXISTING
knobs
331         (``indexing.detector_impl`` / ``skip_ai_handoff``, ``storage.backend``) in the
loader.
332
333         Valid values mirror ``backend/config/presets.py`` (parity-pinned by a unit test)."""
334
335         # "" = no profile (today's manual knobs). Others apply presets.PROFILE_PRESETS.
336         profile: Literal["", "truly-local", "cloud-serving", "cpu-mock"] = ""
337         # Where the core runs. Set by the profile preset; overridable in config.json. Inert
in M1
338         # (recorded for observability; consumed by the ComputeBackend seam from M2+).
339         compute_target: Literal["", "local_gpu", "cloud_run", "cpu_mock"] = ""
340
341
342     class AppConfig(BaseModel):
343         """Root configuration object.
344
345         All runtime configuration should be accessed via an instance of this class
346         (or the loader) rather than raw dict lookups.
347         """
348
349         sport: Sport = "badminton"
350         ai_provider: str = "gemini"
351         # Auto-tracking alias: 2.5-flash is legacy now that 3.5-flash is the stable
flagship;

```

```

352         # the alias never deprecates underneath us. Eval reports record the run date
alongside.
353         ai_model: str = "gemini-flash-latest"
354
355         validation: ValidationConfig = Field(default_factory=ValidationConfig)
356         indexing: IndexingConfig = Field(default_factory=IndexingConfig)
357         output: OutputConfig = Field(default_factory=OutputConfig)
358         stitching: StitchingConfig = Field(default_factory=StitchingConfig)
359         security: SecurityConfig = Field(default_factory=SecurityConfig)
360         logging: LoggingConfig = Field(default_factory=LoggingConfig)
361         storage: StorageConfig = Field(default_factory=StorageConfig)
362         deployment: DeploymentConfig = Field(default_factory=DeploymentConfig)
363
364         # Allow extra keys during transition so existing config.json files do not break.
365         model_config = {
366             "extra": "allow",
367             "validate_assignment": True,
368         }
369
370         # Temporary migration shim for legacy dict-style access (config.get("section",
{ }).get(...)).
371         # Will be removed once all call sites (validators, segmenters, etc.) move to typed
372         # attribute access / the typed config. A nested *section* is returned as a plain
dict
373         # via model_dump, so legacy chained access keeps working ? returning the Pydantic
374         # section model instead breaks it (a model has no .get), which bit every validator
on
375         # real runs. Scalar leaves are returned as-is.
376         def get(self, key: str, default: Any = None) -> Any:
377             _MISSING = object()
378             val: Any = _MISSING
379             if key in type(self).model_fields or hasattr(self, key):
380                 val = getattr(self, key, _MISSING)
381             if val is _MISSING:
382                 # Legacy convenience: find a leaf key inside any nested section. Sections
are
383                 # discovered dynamically from the model fields, so a future top-level
section
384                 # (e.g. `audio`, `jobs`) is covered automatically without editing a
hardcoded list.
385                 for fname in type(self).model_fields:
386                     section = getattr(self, fname, None)
387                     if isinstance(section, BaseModel) and hasattr(section, key):
388                         val = getattr(section, key)
389                         break
390             if val is _MISSING:
391                 return default
392             if isinstance(val, BaseModel):
393                 return val.model_dump(mode="json")
394             return val
395
396         def __getitem__(self, key: str) -> Any:
397             val = self.get(key)
398             if val is None and key not in type(self).model_fields:
399                 # Re-raise like a real dict for unknown top-level
400                 raise KeyError(key)
401             return val
402
403
404         # Convenience type alias for callers that want to annotate
405         Config = AppConfig
406

```

```

1      """Configuration loader and saver.
2
3      Provides a single place to load validated AppConfig with built-in defaults
4      merged with optional file overrides (config.json).
5
6      This replaces scattered json.load + .get() patterns throughout the codebase.
7      """
8
9      from __future__ import annotations
10
11     import json
12     import logging
13     import os
14     from pathlib import Path
15     from typing import Any
16
17     from pydantic import BaseModel
18
19     from .models import AppConfig
20     from .presets import (
21         COMPUTE_TARGET_FOR_PROFILE,
22         PROFILE_PRESETS,
23         is_known_profile,
24         preset_for,
25         suggest_profile,
26     )
27
28     logger = logging.getLogger(__name__)
29
30     DEFAULT_CONFIG_PATH = Path("config.json")
31
32
33     def _deep_merge(base: dict[str, Any], override: dict[str, Any]) -> dict[str, Any]:
34         """Recursively overlay ``override`` onto ``base`` (override wins; dicts merge,
35         scalars
36         replace). Used ONLY for the profile-preset precedence path so a config.json sub-key
37         overrides a preset sub-key while sibling preset/default sub-keys survive ? the no-
38         profile
39         path keeps the original shallow ``{**defaults, **file_data}`` merge unchanged."""
40         out = dict(base)
41         for key, val in override.items():
42             if isinstance(val, dict) and isinstance(out.get(key), dict):
43                 out[key] = _deep_merge(out[key], val)
44             else:
45                 out[key] = val
46         return out
47
48     def _resolve_explicit_profile(file_data: dict[str, Any]) -> str:
49         """The EXPLICITLY chosen deployment profile, env winning over config.json. ``""`` if
50         none
51         is set. Auto-detect is deliberately NOT consulted here ? it only suggests (D15), so
52         a bare
53         environment never silently selects a profile (and never silently spends on
54         cloud)."""
55         env_profile = os.environ.get("DEPLOYMENT_PROFILE")
56         if env_profile and env_profile.strip():
57             return env_profile.strip()
58         dep = file_data.get("deployment")
59         if isinstance(dep, dict):
60             prof = dep.get("profile")
61             if isinstance(prof, str) and prof.strip():
62                 return prof.strip()
63         return ""

```

[illegible]

```

119         file_data = {}
120
121     if file_data:
122         unknown = _unknown_key_paths(file_data, AppConfig)
123         if unknown:
124             logger.warning(
125                 "[CONFIG WARNING] %s contains keys the typed config does not "
126                 "recognize (typo'd or removed knobs ? top-level extras are kept "
127                 "but unread; unknown section keys are silently dropped): %s",
128                 cfg_path, ", ".join(sorted(unknown)),
129             )
130
131     # Precedence: built-in defaults < profile preset < config.json (< env/--set, applied
at
132     # consumption sites downstream). A *deployment profile* is opt-in ? it is applied
ONLY when
133     # explicitly selected (DEPLOYMENT_PROFILE env or config.json deployment.profile).
With no
134     # profile the merge is byte-identical to the pre-M1 loader, so the default path is
unchanged.
135     defaults = _get_builtin_defaults()
136     profile = _resolve_explicit_profile(file_data)
137
138     if profile and not is_known_profile(profile):
139         # An unrecognised EXPLICIT profile is loud, not silent (the loader's "dropped
config"
140         # ethos): warn and apply no preset. A bad value coming from config.json
additionally
141         # hits the typed Literal at model_validate below (a hard, clear error); a bad
value
142         # coming only from the env override is ignored here with this warning.
143         logger.warning(
144             "[CONFIG] unrecognized deployment profile %r (valid: %s); applying no
preset.",
145             profile, ", ".join(PROFILE_PRESETS),
146         )
147         profile = ""
148
149     if profile:
150         merged = _deep_merge(_deep_merge(defaults, preset_for(profile)), file_data)
151         # Record the profile actually applied ? the env may have chosen it over
config.json,
152         # so re-stamp it onto the merged deployment section to keep the record honest.
153         merged.setdefault("deployment", {})
154         if isinstance(merged["deployment"], dict):
155             merged["deployment"]["profile"] = profile
156             ct = merged["deployment"].get("compute_target")
157             canonical = COMPUTE_TARGET_FOR_PROFILE.get(profile)
158             if ct and canonical and ct != canonical:
159                 logger.warning(
160                     "[CONFIG] deployment profile %r normally uses compute_target=%r but
"
161                     "config.json overrides it to %r ? honoring the explicit override.",
162                     profile, canonical, ct,
163                 )
164             logger.info("[CONFIG] deployment profile ACTIVE: %s (compute_target=%s)",
165                         profile, merged["deployment"].get("compute_target"))
166     else:
167         # No profile ? the original shallow merge (Pydantic re-defaults nested
sections).
168         merged = {**defaults, **file_data}
169         suggestion = suggest_profile() # cheap, env-only (no torch); SUGGESTS only
(D15)
170         if suggestion:
171             logger.debug(

```

```

172             "[CONFIG] no deployment.profile set (manual knobs). Environment suggests
%r ? "
173             "set deployment.profile or DEPLOYMENT_PROFILE to apply (auto-detect
never "
174             "auto-applies ? D15).", suggestion,
175             )
176
177     return AppConfig.model_validate(merged)
178
179
180 def save_default_config(path: str | Path | None = None) -> Path:
181     """Write a fresh default config.json.
182
183     Used by scripts/doctor.py diagnostics and the --setup CLI path.
184     Overwrites if the file exists.
185     """
186     cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
187     cfg = AppConfig()
188     cfg_path.parent.mkdir(parents=True, exist_ok=True)
189     cfg_path.write_text(cfg.model_dump_json(indent=4), encoding="utf-8")
190     return cfg_path
191
192
193 # Backwards-compatible alias used by some older call sites during transition.
194 def get_default_config() -> dict[str, Any]:
195     """Return defaults as a plain dict (for transitional use only)."""
196     return _get_builtin_defaults()
197

```

9. user (2026-06-21T07:24:08.841Z)

```

1     """Configuration package.
2
3     Public API:
4         from backend.config import load_config, AppConfig, Config
5
6     This package is the canonical place for all configuration access.
7     """
8
9     from .loader import load_config, save_default_config
10    from .models import ( # re-export common ones
11        AppConfig,
12        Config,
13        DeploymentConfig,
14        IndexingConfig,
15        StitchingConfig,
16        StorageConfig,
17        ValidationConfig,
18    )
19    from .presets import PROFILE_PRESETS, suggest_profile
20
21    __all__ = [
22        "load_config",
23        "save_default_config",
24        "AppConfig",
25        "Config",
26        "DeploymentConfig",
27        "IndexingConfig",
28        "StitchingConfig",
29        "StorageConfig",
30        "ValidationConfig",
31        "PROFILE_PRESETS",
32        "suggest_profile",
33    ]
34

```

10. user (2026-06-21T07:24:09.254Z)

```
1      """Tests for the deployment-profile config (COMPUTE_DECOUPLED_SERVING M1).
2
3      What this pins:
4      - The default profile ("") applies NO preset ? the loaded config is byte-identical to
the
5      pre-M1 default, so the core is unaffected (the project's "default-OFF, zero core
change").
6      - Each profile applies its coherent preset bundle (compute_target + the existing
7      detector_impl / skip_ai_handoff / storage.backend knobs).
8      - Precedence: built-in defaults < profile preset < config.json (a file sub-key overrides
a
9      preset sub-key while sibling preset/default sub-keys survive ? i.e. real deep-merge
layering).
10     - D15: auto-detect SUGGESTS but never auto-applies (a bare cloud-like env does NOT
silently
11     select a profile / spend on cloud).
12     - The model's Literal enum stays in parity with backend/config/presets.PROFILE_PRESETS.
13     """
14
15     import json
16     import logging
17     import os
18
19     import pytest
20     from pydantic import ValidationError
21
22     from backend.config import AppConfig, DeploymentConfig, load_config, save_default_config
23     from backend.config.presets import (
24         COMPUTE_TARGET_FOR_PROFILE,
25         COMPUTE_TARGETS,
26         PROFILE_PRESETS,
27         PROFILES,
28         is_known_profile,
29         preset_for,
30         suggest_profile,
31     )
32
33     NONEXISTENT = "/nonexistent/path/12345/deployment_config.json"
34
35
36     def _write(tmp_path, data) -> str:
37         p = tmp_path / "cfg.json"
38         p.write_text(json.dumps(data))
39         return str(p)
40
41
42     # -----
43     # Defaults / additivity ? the default path must be unchanged.
44     # -----
45     def test_deployment_defaults_are_empty_and_inert():
46         cfg = load_config(NONEXISTENT)
47         assert isinstance(cfg.deployment, DeploymentConfig)
48         assert cfg.deployment.profile == ""
49         assert cfg.deployment.compute_target == ""
50
51
52     def test_default_path_applies_no_preset(monkeypatch):
53         """With no profile selected, the existing knobs keep their built-in defaults (no
preset
54         mutates detector_impl / skip_ai_handoff / storage.backend)."""
55         monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
56         cfg = load_config(NONEXISTENT)
57         assert cfg.indexing.detector_impl == "auto"
```

```

58         assert cfg.indexing.skip_ai_handoff is False
59         assert cfg.storage.backend == "local"
60
61
62     def test_default_config_dump_only_adds_deployment_section():
63         """Adding the section must not perturb any other default ? the dump equals a fresh
64         AppConfig()'s dump, and the new section is present with empty values."""
65         dumped = load_config(NONEXISTENT).model_dump(mode="json")
66         assert dumped == AppConfig().model_dump(mode="json")
67         assert dumped["deployment"] == {"profile": "", "compute_target": ""}
68
69
70     # -----
71     # Parity: the model Literal must match the presets table (no drift).
72     # -----
73     def test_profile_literal_matches_presets():
74         import typing
75
76         literal_profiles =
77 set(typing.get_args(DeploymentConfig.model_fields["profile"].annotation))
78         assert literal_profiles == {"", *PROFILES}
79         assert set(PROFILE_PRESETS) == set(PROFILES)
80         assert set(COMPUTE_TARGET_FOR_PROFILE) == set(PROFILES)
81
82     def test_compute_target_literal_matches_presets():
83         import typing
84
85         literal_targets =
86 set(typing.get_args(DeploymentConfig.model_fields["compute_target"].annotation))
87         assert literal_targets == {"", *COMPUTE_TARGETS}
88         assert set(COMPUTE_TARGET_FOR_PROFILE.values()) == set(COMPUTE_TARGETS)
89
90     def test_each_preset_only_touches_existing_knobs():
91         """A preset must only set knobs that exist on the typed models ? otherwise it would
92 be a
93 silent no-op. Validates the merged result against AppConfig for every profile."""
94         for profile in PROFILES:
95             # A preset's value set must validate cleanly when merged onto defaults.
96             merged = {**AppConfig().model_dump(mode="json")}
97             for section, kv in preset_for(profile).items():
98                 merged.setdefault(section, {})
99                 if isinstance(merged[section], dict):
100                     merged[section].update(kv)
101             AppConfig.model_validate(merged) # raises if a preset key/value is invalid
102
103     # -----
104     # Per-profile preset application.
105     # -----
106     def test_cpu_mock_profile(tmp_path):
107         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "cpu-mock"}}))
108         assert cfg.deployment.profile == "cpu-mock"
109         assert cfg.deployment.compute_target == "cpu_mock"
110         assert cfg.indexing.detector_impl == "stub"
111         assert cfg.indexing.skip_ai_handoff is True
112         assert cfg.storage.backend == "local"
113
114
115     def test_cloud_serving_profile(tmp_path):
116         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "cloud-serving"}}))
117         assert cfg.deployment.profile == "cloud-serving"
118         assert cfg.deployment.compute_target == "cloud_run"
119         assert cfg.indexing.detector_impl == "native"

```

```

120         assert cfg.indexing.skip_ai_handoff is False # Gemini-ON until the owned model
clears the floor
121         assert cfg.storage.backend == "gcs"
122
123
124     def test_truly_local_profile(tmp_path):
125         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "truly-local"}}))
126         assert cfg.deployment.profile == "truly-local"
127         assert cfg.deployment.compute_target == "local_gpu"
128         # detector_impl is intentionally left at the production default (env picks WSL vs
native).
129         assert cfg.indexing.detector_impl == "auto"
130         assert cfg.indexing.skip_ai_handoff is True # zero-cloud
131         assert cfg.storage.backend == "local"
132
133
134     # -----
135     # Precedence: defaults < preset < config.json (deep-merge layering).
136     # -----
137     def test_config_json_overrides_preset_subkey(tmp_path):
138         """A file sub-key beats the preset's sub-key, but a sibling preset sub-key in the
SAME
139         section survives (this is the deep-merge layering ? a shallow merge would wipe
it)."""
140         cfg = load_config(_write(tmp_path, {
141             "deployment": {"profile": "cpu-mock"},
142             "indexing": {"detector_impl": "native"}, # override one preset knob...
143         }))
144         assert cfg.indexing.detector_impl == "native" # file wins
145         assert cfg.indexing.skip_ai_handoff is True # ...preset sibling SURVIVES (deep
merge)
146
147
148     def test_preset_does_not_wipe_unrelated_defaults(tmp_path):
149         """Selecting a profile + setting an unrelated knob keeps the rest of the section at
its
150         built-in default (the preset and the file both layer onto the full defaults)."""
151         cfg = load_config(_write(tmp_path, {
152             "deployment": {"profile": "cpu-mock"},
153             "indexing": {"ai_max_workers": 2},
154         }))
155         assert cfg.indexing.detector_impl == "stub" # preset
156         assert cfg.indexing.skip_ai_handoff is True # preset
157         assert cfg.indexing.ai_max_workers == 2 # file
158         assert cfg.indexing.default_segmenter == "yolo_hybrid" # untouched default
159         assert cfg.indexing.max_window_duration == 6.0 # untouched default
160
161
162     def test_config_json_can_override_compute_target_with_warning(tmp_path, caplog):
163         cfg_path = _write(tmp_path, {
164             "deployment": {"profile": "cpu-mock", "compute_target": "cloud_run"},
165         })
166         with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
167             cfg = load_config(cfg_path)
168             assert cfg.deployment.profile == "cpu-mock"
169             assert cfg.deployment.compute_target == "cloud_run" # explicit override honored
(D15)
170             assert any("overrides it to 'cloud_run'" in r.message for r in caplog.records)
171
172
173     # -----
174     # Env selection + precedence over config.json.
175     # -----
176     def test_env_selects_profile(tmp_path, monkeypatch):
177         monkeypatch.setenv("DEPLOYMENT_PROFILE", "cpu-mock")

```

```

178     cfg = load_config(_write(tmp_path, {}))
179     assert cfg.deployment.profile == "cpu-mock"
180     assert cfg.indexing.detector_impl == "stub"
181
182
183     def test_env_profile_overrides_config_json_profile(tmp_path, monkeypatch):
184         """env wins over config.json for the profile CHOICE, and the recorded profile
matches the
185         preset actually applied (not the file's value)."""
186         monkeypatch.setenv("DEPLOYMENT_PROFILE", "cpu-mock")
187         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "cloud-serving"}}))
188         assert cfg.deployment.profile == "cpu-mock"          # env value recorded
189         assert cfg.deployment.compute_target == "cpu_mock"   # cpu-mock preset applied
190         assert cfg.indexing.detector_impl == "stub"
191
192
193     def test_blank_env_profile_is_ignored(tmp_path, monkeypatch):
194         monkeypatch.setenv("DEPLOYMENT_PROFILE", " ")
195         cfg = load_config(_write(tmp_path, {}))
196         assert cfg.deployment.profile == ""
197         assert cfg.indexing.detector_impl == "auto"
198
199
200     def test_unknown_env_profile_warns_and_applies_no_preset(tmp_path, monkeypatch, caplog):
201         monkeypatch.setenv("DEPLOYMENT_PROFILE", "banana")
202         with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
203             cfg = load_config(_write(tmp_path, {}))
204             assert cfg.deployment.profile == ""              # not applied
205             assert cfg.indexing.detector_impl == "auto"      # no preset
206             assert any("unrecognized deployment profile 'banana'" in r.message for r in
caplog.records)
207
208
209     def test_unknown_profile_in_config_json_raises(tmp_path, monkeypatch):
210         """A bad VALUE in config.json hits the typed Literal ? a loud, hard error (same
class as a
211         bad `sport`), not a silent drop."""
212         monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
213         with pytest.raises(ValidationError):
214             load_config(_write(tmp_path, {"deployment": {"profile": "banana"}}))
215
216
217     # -----
218     # D15 ? auto-detect SUGGESTS but never AUTO-APPLIES.
219     # -----
220     def test_cloud_env_does_not_auto_apply_profile(tmp_path, monkeypatch):
221         """K_SERVICE (Cloud Run) makes suggest_profile() recommend cloud-serving, but with
no
222         explicit profile the loaded config stays fully default ? NO silent cloud
selection."""
223         monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
224         monkeypatch.setenv("K_SERVICE", "khehsutra-serving")
225         cfg = load_config(_write(tmp_path, {}))
226         assert cfg.deployment.profile == ""                  # NOT auto-applied
227         assert cfg.deployment.compute_target == ""
228         assert cfg.indexing.detector_impl == "auto"          # default, not "native"
229         assert cfg.storage.backend == "local"                 # default, not "gcs"
230
231
232     def test_active_profile_is_logged(tmp_path, caplog):
233         with caplog.at_level(logging.INFO, logger="backend.config.loader"):
234             load_config(_write(tmp_path, {"deployment": {"profile": "truly-local"}}))
235             assert any("deployment profile ACTIVE: truly-local" in r.message for r in
caplog.records)
236

```

```

237
238 def test_no_active_profile_info_when_unset(tmp_path, caplog, monkeypatch):
239     monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
240     with caplog.at_level(logging.INFO, logger="backend.config.loader"):
241         load_config(_write(tmp_path, {}))
242     assert not any("profile ACTIVE" in r.message for r in caplog.records)
243
244
245 # -----
246 # suggest_profile() ? pure env-only logic + the opt-in cuda probe.
247 # -----
248 def test_suggest_profile_cloud_run():
249     assert suggest_profile({"K_SERVICE": "svc"}) == "cloud-serving"
250
251
252 def test_suggest_profile_ci():
253     assert suggest_profile({"CI": "true"}) == "cpu-mock"
254     assert suggest_profile({"CI": "1"}) == "cpu-mock"
255     assert suggest_profile({"CI": "yes"}) == "cpu-mock"
256
257
258 def test_suggest_profile_ci_falsy_or_absent():
259     assert suggest_profile({}) == ""
260     assert suggest_profile({"CI": "false"}) == ""
261     assert suggest_profile({"CI": "0"}) == ""
262     assert suggest_profile({"CI": ""}) == ""
263
264
265 def test_suggest_profile_k_service_beats_ci():
266     assert suggest_profile({"K_SERVICE": "svc", "CI": "true"}) == "cloud-serving"
267
268
269 def test_suggest_profile_no_cuda_probe_by_default(monkeypatch):
270     """probe_cuda defaults False ? no torch import, no truly-local suggestion even with
a GPU."""
271     import backend.config.presets as presets
272
273     monkeypatch.setattr(presets, "_cuda_available", lambda: True)
274     assert suggest_profile({}) == "" # probe off ? no suggestion
275     assert suggest_profile({}, probe_cuda=True) == "truly-local"
276
277
278 def test_suggest_profile_cuda_probe_false(monkeypatch):
279     import backend.config.presets as presets
280
281     monkeypatch.setattr(presets, "_cuda_available", lambda: False)
282     assert suggest_profile({}, probe_cuda=True) == ""
283
284
285 def test_cuda_available_never_raises(monkeypatch):
286     """Even a broken/missing torch must degrade to False, never propagate."""
287     import builtins
288
289     import backend.config.presets as presets
290
291     real_import = builtins.__import__
292
293     def boom(name, *a, **k):
294         if name == "torch":
295             raise RuntimeError("no torch")
296         return real_import(name, *a, **k)
297
298     monkeypatch.setattr(builtins, "__import__", boom)
299     assert presets._cuda_available() is False
300

```

```

301
302 # -----
303 # Helpers + plumbing.
304 # -----
305 def test_is_known_profile():
306     assert is_known_profile("cpu-mock") is True
307     assert is_known_profile("") is False
308     assert is_known_profile("banana") is False
309
310
311 def test_preset_for_returns_copy():
312     a = preset_for("cpu-mock")
313     a["indexing"]["detector_impl"] = "mutated"
314     b = preset_for("cpu-mock")
315     assert b["indexing"]["detector_impl"] == "stub" # module table not mutated
316     assert preset_for("") == {}
317     assert preset_for("banana") == {}
318
319
320 def test_save_default_config_includes_deployment(tmp_path):
321     p = save_default_config(str(tmp_path / "out.json"))
322     saved = json.loads(p.read_text())
323     assert saved["deployment"] == {"profile": "", "compute_target": ""}
324     cfg = load_config(str(p))
325     assert cfg.deployment.profile == ""
326
327
328 def test_repo_config_json_has_no_profile():
329     """The shipped config.json must not silently opt into a profile (default path stays
today's
330     behaviour for the real deployment)."""
331     repo_root = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
332     cfg_path = os.path.join(repo_root, "config.json")
333     if not os.path.exists(cfg_path):
334         pytest.skip("no repo config.json")
335     cfg = load_config(cfg_path)
336     assert cfg.deployment.profile == ""
337

```

11. user (2026-06-21T07:24:09.703Z)

```

1     """Basic tests for the typed configuration package.
2
3     These ensure:
4     - Defaults are sensible and come from the models.
5     - File overrides work.
6     - Old config files with extra keys are tolerated.
7     - save_default_config produces loadable output.
8     """
9     import json
10    import tempfile
11    import os
12
13
14    from backend.config import load_config, save_default_config, AppConfig
15
16
17    def test_load_default_config():
18        # Force load without the project's config.json so we test the model built-in
defaults
19        cfg = load_config("/nonexistent/path/that/does/not/exist/12345/config.json")
20        assert isinstance(cfg, AppConfig)
21        assert cfg.sport == "badminton"
22        assert cfg.indexing.default_segmenter == "yolo_hybrid"
23        assert cfg.validation.min_blur_threshold == 20.0 # updated default from review

```

```

24         assert cfg.indexing.ai_max_workers == 5
25
26
27     def test_logging_config_default_quiets_http():
28         """verbose_http defaults to False so run.log is quiet unless explicitly turned
29 up."""
29         cfg = load_config("/nonexistent/path/12345/config.json")
30         assert cfg.logging.verbose_http is False
31
32
33     def test_skip_ai_handoff_default_off():
34         """Fully-local mode is opt-in: skip_ai_handoff defaults False so the AI handoff
35 runs."""
35         cfg = load_config("/nonexistent/path/12345/config.json")
36         assert cfg.indexing.skip_ai_handoff is False
37
38
39     def test_skip_ai_handoff_file_override():
40         with tempfile.TemporaryDirectory() as td:
41             p = os.path.join(td, "skip.json")
42             with open(p, "w") as fh:
43                 json.dump({"indexing": {"skip_ai_handoff": True}}, fh)
44             cfg = load_config(p)
45             assert cfg.indexing.skip_ai_handoff is True
46
47
48     def test_logging_config_file_override():
49         with tempfile.TemporaryDirectory() as td:
50             p = os.path.join(td, "logging.json")
51             with open(p, "w") as fh:
52                 json.dump({"logging": {"verbose_http": True}}, fh)
53             cfg = load_config(p)
54             assert cfg.logging.verbose_http is True
55
56
57     def test_fusion_config_defaults_reproduce_control():
58         """The fusion block defaults must be no-ops: wasb substrate, no cues, no gating, no
59 mask.
60 Gate A (min_crossings) defaults to 0 = OFF: the #146 served-default flip was
61 REVERTED
62 (eval?serve ? the +0.074 was on the harness's recall-oriented PROPOSAL windowing; on
63 the
64 served COARSE windowing Gate A is neutral/negative, ?0.008, dropping real rallies).
65 See
66 docs/QUALITY_ITERATIONS.md and the served litmus in
67 tests/test_served_gate_a_regression.py.
68 """
69         cfg = load_config("/nonexistent/path/12345/config.json")
70         f = cfg.indexing.fusion
71         assert f.substrate == "wasb"
72         assert f.cue_flags == {} # person cue OFF
73         assert f.min_crossings == 0 # Gate A OFF (default) ? #146 served flip
74         reverted
75         assert f.min_players == 0 # Gate B OFF
76         assert f.court_polygon is None # no mask
77         assert f.velocity_threshold == 0.02
78
79
80     def test_fusion_config_gate_a_opt_in_via_min_crossings():
81         """Gate A is opt-in: setting min_crossings explicitly enables it (the eval-only
82 operating
83 point), while the served default stays OFF (#146 flip reverted)."""
84         with tempfile.TemporaryDirectory() as td:
85             p = os.path.join(td, "optin.json")

```

```

80         with open(p, "w") as fh:
81             json.dump({"indexing": {"fusion": {"min_crossings": 2}}}, fh)
82         cfg = load_config(p)
83         assert cfg.indexing.fusion.min_crossings == 2    # Gate A ON only when explicitly
set
84
85
86     def test_served_default_segmenter_is_gemini():
87         """The shipped config.json served default is the prior detector (gemini) ? the #146
flip
88         to fusion_hybrid was REVERTED as a confirmed served regression (eval?serve)."""
89         repo_root = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
90         cfg = load_config(os.path.join(repo_root, "config.json"))
91         assert cfg.indexing.default_segmenter == "gemini"
92         # ...and the served fusion default keeps Gate A OFF (no served gating).
93         assert cfg.indexing.fusion.min_crossings == 0
94
95
96     def test_fusion_config_file_override():
97         with tempfile.TemporaryDirectory() as td:
98             p = os.path.join(td, "fusion.json")
99             with open(p, "w") as fh:
100                 json.dump({"indexing": {"fusion": {
101                     "substrate": "tracknet", "min_crossings": 2,
102                     "cue_flags": {"person": True}, "court_polygon": [[0.1, 0.1], [0.9, 0.1],
[0.5, 0.9]],
103                 }}, fh)
104             cfg = load_config(p)
105             f = cfg.indexing.fusion
106             assert f.substrate == "tracknet"
107             assert f.min_crossings == 2
108             assert f.cue_flags == {"person": True}
109             assert len(f.court_polygon) == 3
110             # An old config without a fusion block still loads (defaults applied).
111             assert cfg.indexing.wasb.velocity_threshold == 0.02
112
113
114     def test_file_override():
115         with tempfile.TemporaryDirectory() as td:
116             p = os.path.join(td, "override.json")
117             data = {
118                 "indexing": {
119                     "default_segmenter": "gemini",
120                     "ai_max_workers": 3,
121                 },
122                 "validation": {
123                     "min_blur_threshold": 42.0,
124                 },
125             }
126             with open(p, "w") as f:
127                 json.dump(data, f)
128
129             cfg = load_config(p)
130             assert cfg.indexing.default_segmenter == "gemini"
131             assert cfg.indexing.ai_max_workers == 3
132             assert cfg.validation.min_blur_threshold == 42.0
133             # Unspecified keys fall back to defaults
134             assert cfg.sport == "badminton"
135
136
137     def test_extra_keys_tolerated():
138         """Old or future keys in config.json must not break loading."""
139         with tempfile.TemporaryDirectory() as td:
140             p = os.path.join(td, "with_extra.json")
141             data = {

```

```

142         "sport": "badminton",
143         "indexing": {"default_segmenter": "yolo_hybrid"},
144         "some_future_key": "ignored_for_now",
145         "validation": {"relevance": {"court_pixels_min_ratio": 0.1}},
146     }
147     with open(p, "w") as f:
148         json.dump(data, f)
149
150     cfg = load_config(p)
151     assert cfg.indexing.default_segmenter == "yolo_hybrid"
152     assert cfg.validation.relevance.court_pixels_min_ratio == 0.1
153
154
155 def test_unknown_keys_warn_with_dotted_paths(tmp_path, caplog):
156     """Typo'd keys must be LOUD: top-level extras are kept-but-unread and unknown
157     section keys are silently dropped, so without this warning a typo just makes
158     the knob a no-op (the audit's 'silently dropped config' class)."""
159     p = tmp_path / "typo.json"
160     p.write_text(json.dumps({
161         "some_future_key": 1,                                # top-level extra
162         "indexing": {"velocty_threshold": 0.5},              # typo'd section key
163         "validation": {"relevance": {"unknown_gate": True}}, # doubly nested
164     }))
165     import logging
166     with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
167         cfg = load_config(str(p))
168         assert isinstance(cfg, AppConfig) # still non-fatal
169         joined = " ".join(r.message for r in caplog.records)
170         assert "some_future_key" in joined
171         assert "indexing.velocty_threshold" in joined
172         assert "validation.relevance.unknown_gate" in joined
173
174
175 def test_recognized_config_does_not_warn(tmp_path, caplog):
176     p = tmp_path / "clean.json"
177     p.write_text(json.dumps({
178         "sport": "badminton",
179         "indexing": {"default_segmenter": "gemini", "wasb": {"velocity_threshold":
0.02}},
180     }))
181     import logging
182     with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
183         load_config(str(p))
184     assert not [r for r in caplog.records if "CONFIG WARNING" in r.message]
185
186
187 def test_non_object_config_degrades_to_defaults(tmp_path, caplog):
188     """A config.json that is valid JSON but NOT an object (`[]`, `[1]`, `\"x\"`, `42`,
189     `true`,
190     `null`) must degrade to a fully-defaulted AppConfig + a warning ? never crash
191     startup.
192     Regression: previously a truthy non-mapping hit `.items()` (AttributeError) and a
193     falsy one
194     broke the `{**defaults, **file_data}` unpack (TypeError), turning an operator typo
195     into a hard
196     server/CLI startup crash."""
197     import logging
198     default_dump = AppConfig().model_dump(mode="json")
199     for raw in ("[]", "[1, 2]", "\"x\"", "42", "true", "null"):
200         p = tmp_path / "bad.json"
201         p.write_text(raw)
202         caplog.clear()
203         with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
204             cfg = load_config(str(p)) # must not raise
205         assert isinstance(cfg, AppConfig)

```

```

202         assert cfg.model_dump(mode="json") == default_dump, f"{raw} should yield pure
defaults"
203         assert any("not a JSON object" in r.message for r in caplog.records), \
204             f"{raw} should log the non-object warning"
205
206
207     def test_save_default_config_roundtrip():
208         with tempfile.TemporaryDirectory() as td:
209             p = os.path.join(td, "saved.json")
210             saved = save_default_config(p)
211             assert str(saved) == p
212             assert os.path.exists(p)
213
214             cfg = load_config(p)
215             assert cfg.indexing.default_segmenter == "yolo_hybrid"
216             assert cfg.validation.min_blur_threshold == 20.0
217
218
219     def test_model_is_the_source_of_truth():
220         """Changing a default in the model should be reflected without touching json."""
221         # This is a documentation test more than anything.
222         cfg = AppConfig()
223         assert cfg.indexing.tracknet.velocity_threshold == 0.02
224
225
226     def test_output_dir_field_default_and_override():
227         """output_dir now lives on OutputConfig (was a write-only runtime hack), so
228         /api/compile and the CLI resolve the same value. Default + file override + the
229         runtime mutation that backend/main.py::main performs for --output-dir."""
230         cfg = AppConfig()
231         assert cfg.output.output_dir == "output" # built-in default
232
233         # The CLI override path: main() assigns args.output_dir onto the typed model.
234         cfg.output.output_dir = "/tmp/custom_out"
235         assert cfg.output.output_dir == "/tmp/custom_out"
236
237         # File override is honored on load.
238         with tempfile.TemporaryDirectory() as td:
239             p = os.path.join(td, "out_override.json")
240             with open(p, "w") as f:
241                 json.dump({"output": {"output_dir": "reels", "db_name": "x.db"}}, f)
242             loaded = load_config(p)
243             assert loaded.output.output_dir == "reels"
244             assert loaded.output.db_name == "x.db"
245
246     def test_config_shallow_section_replacement():
247         """Positive test for Item 2 (config loader): shallow section replacement.
248
249         File-provided 'indexing' section replaces the defaults 'indexing' section entirely
250         (no deep merge of sub-keys), per loader's {**defaults, **file_data}.
251         """
252         with tempfile.TemporaryDirectory() as td:
253             p = os.path.join(td, "shallow.json")
254             with open(p, "w") as fh:
255                 json.dump({"indexing": {"default_segmenter": "motion"}}, fh)
256             cfg = load_config(p)
257             assert cfg.indexing.default_segmenter == "motion"
258             # Pydantic defaults fill the rest of the (replaced) section
259             assert cfg.indexing.ai_max_workers == 5
260
261
262     def test_config_unknown_key_paths_warning_and_tolerate():
263         """Positive test for Item 2: _unknown_key_paths warning for unknown top-level keys,
264         but tolerated (extra='allow'), and unknown section keys dropped silently.
265         """

```

```

266     with tempfile.TemporaryDirectory() as td:
267         p = os.path.join(td, "unknown.json")
268         with open(p, "w") as fh:
269             json.dump({
270                 "indexing": {"default_segmenter": "motion"},
271                 "completely_unknown_top": "bar",
272                 "validation": {"unknown_in_section": 99} # dropped
273             }, fh)
274         cfg = load_config(p)
275         assert cfg.indexing.default_segmenter == "motion"
276         # tolerated, no crash
277
278
279     def test_get_returns_dict_for_nested_sections_legacy_compat():
280         """Regression: legacy consumers (validators/segmenters) do
281         ``config.get("validation", {}).get("min_resolution_width", ...)``. The dict-compat
282         .get must return nested SECTIONS as plain dicts (not Pydantic models), or that
283         chained access raises 'X object has no attribute get' on every real run.
284
285         This is the original strong legacy test (restored per #160 review). It covers the
286         full set of footgun guards: dict sections, nested dicts, scalar leaves with
defaults,
287         missing key fallbacks, and leaf keys discoverable inside sections.
288         """
289         cfg = load_config("/nonexistent/123/config.json")
290         validation = cfg.get("validation", {})
291         assert isinstance(validation, dict)
292         assert validation.get("min_resolution_width") == 1280
293         # nested-of-nested is also dict-accessible
294         assert isinstance(validation.get("relevance"), dict)
295         assert validation["relevance"].get("court_pixels_min_ratio") == 0.05
296
297         # scalar leaves are returned unchanged; a present scalar ignores the default arg;
298         # a missing key falls back to the default
299         assert cfg.get("ai_provider") == "gemini"
300         assert cfg.get("ai_provider", "ignored-default") == "gemini"
301         assert cfg.get("does_not_exist", "fallback") == "fallback"
302
303         # leaf keys inside any section resolve via the dynamic section scan (not a hardcoded
304         # list), so a future top-level section's keys would be discoverable too
305         assert cfg.get("min_resolution_width") == 1280 # lives inside `validation`
306         assert cfg.get("begin_padding") == 0.5 # lives inside `stitching`
307
308         # legacy chained access still works
309         assert cfg.get("indexing", {}).get("default_segmenter") == "yolo_hybrid"
310
311
312     def test_config_roundtrip_and_typed_reaches_pipeline(monkeypatch):
313         """Positive for Item 2: save_default_config + load_config round-trips;
314         typed AppConfig is what reaches pipeline code (e.g. main uses global_config as
AppConfig).
315         """
316         with tempfile.TemporaryDirectory() as td:
317             p = os.path.join(td, "roundtrip.json")
318             save_default_config(p)
319             cfg = load_config(p)
320             assert isinstance(cfg, AppConfig)
321             assert cfg.sport == "badminton"
322
323         # positive that typed is used in pipeline code (global_config is set at startup in
main());
324         # monkeypatch here so the test does not depend on server initialization)
325         import backend.main as main_module
326         from backend.main import get_config
327         monkeypatch.setattr(main_module, "global_config", cfg)

```

```
328         assert isinstance(get_config(), AppConfig)
329
```

12. assistant (2026-06-21T07:24:13.263Z)

Now let me investigate the backward-compat concern. The critical question is whether adding the `deployment` section to `AppConfig` breaks any caller that enumerates `model\_fields`, dumps config, snapshots it, or asserts an exact field set. Let me grep the repo.

13. user (2026-06-21T07:24:15.726Z)

```
backend\worker\__main__.py:94:     storage_cfg = config.storage.model_dump()
backend\worker\__main__.py:112:         [r.model_dump() for r in val_results])
backend\worker\__main__.py:263:     storage_cfg = config.storage.model_dump()
backend\config\models.py:373:     # via model_dump, so legacy chained access keeps working ?
returning the Pydantic
backend\config\models.py:379:         if key in type(self).model_fields or hasattr(self, key):
backend\config\models.py:385:             for fname in type(self).model_fields:
backend\config\models.py:393:                 return val.model_dump(mode="json")
backend\config\models.py:398:         if val is None and key not in type(self).model_fields:
backend\config\loader.py:71:     fields = model_cls.model_fields
backend\config\loader.py:87:     return AppConfig().model_dump(mode="json")
backend\config\loader.py:189:     cfg_path.write_text(cfg.model_dump_json(indent=4),
encoding="utf-8")
backend\pipeline\stitcher\compiler.py:31:         if hasattr(watermark, "model_dump"):
backend\pipeline\stitcher\compiler.py:32:             return watermark.model_dump()
backend\eval\per_user_regression.py:48:from dataclasses import asdict, dataclass, field
backend\eval\per_user_regression.py:105:     return asdict(self)
backend\eval\per_user_regression.py:129:     return asdict(self)
backend\eval\regression.py:21:from dataclasses import dataclass, asdict
backend\eval\regression.py:63:     return asdict(self)
backend\main.py:87:     return cfg.model_dump(mode="json")
backend\main.py:173:         db_instance.add_video(video_id, req.video_path, "failed",
[r.model_dump() for r in val_results])
backend\main.py:178:         "results": [r.model_dump() for r in val_results],
backend\main.py:184:         db_instance.add_video(video_id, req.video_path, "processed",
[r.model_dump() for r in val_results])
backend\main.py:195:         "results": [r.model_dump() for r in val_results],
backend\main.py:224:         "results": [r.model_dump() for r in val_results],
backend\main.py:237:         "results": [r.model_dump() for r in val_results],
backend\main.py:594:         db_instance.add_video(video_id, args.video_path, status,
[r.model_dump() for r in val_results])
backend\infrastructure\execution_policy.py:22:from dataclasses import asdict, dataclass
backend\infrastructure\execution_policy.py:56:         d = asdict(self)
backend\reporting\__init__.py:45:     if hasattr(config, "model_dump"):
backend\reporting\__init__.py:47:         config = config.model_dump(mode="json")
docs\archives\CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md:46:- `AppConfig.get()`
(models.py:255) is a "temporary migration shim" that returns dict sections (via `model_dump`) or
scalars so legacy chains keep working. It is still load-bearing everywhere.
docs\archives\CODE_AUDIT_AND_TEST_HARDENING-COMPLETED-2026-06-14.md:363: -
`test_config_get_shim_returns_dicts_and_scalars`: exercises the .get() migration shim: returns
dict (via model_dump) for sections, scalar for leaves; supports legacy cfg.get("indexing",
{}).get("default_segmenter").
docs\CODE_MAP.md:59: -> db.add_video(video_id, filepath, status, [r.dict()])
@backend/infrastructure/database.py
docs\CODE_MAP.md:116: - `::load_config` ? ? legacy thin wrapper returning a **dict**
(`load_app_config(path).model_dump(...)`); kept for transition ? new code imports
`load_app_config`/`AppConfig` directly. `::ProcessRequest`/`::QueryRequest`/`::CompileRequest`
pydantic bodies.
docs\CODE_MAP.md:128: - ? `::AppConfig.get` ? ? legacy dict-compat shim: returns nested
**sections as plain dicts** (`model_dump(mode='json')`, NOT the model) so chained
`config.get("indexing", {}).get(...)` keeps working; falls back to searching a leaf key in any
section (sections discovered dynamically from `model_fields`). ? "bit every validator on real
runs" per docstring ? returning the model breaks `.get`. `::AppConfig.__getitem__` raises
`KeyError` for unknown top-level keys else delegates to `.get`.
```

```
docs\CODE_MAP.md:132: - `::DEFAULT_CONFIG_PATH = Path("config.json")` (? CWD-relative, NOT
repo-root). `::get_builtin_defaults` = `AppConfig().model_dump(...)`
docs\CODE_MAP.md:219:[Omitted long matching line]
tests\test_config.py:194:     default_dump = AppConfig().model_dump(mode="json")
tests\test_config.py:202:     assert cfg.model_dump(mode="json") == default_dump, f"{raw}
should yield pure defaults"
tests\test_config_deployment.py:65:     dumped = load_config(NONEXISTENT).model_dump(mode="json")
tests\test_config_deployment.py:66:     assert dumped == AppConfig().model_dump(mode="json")
tests\test_config_deployment.py:76:     literal_profiles =
set(typing.get_args(DeploymentConfig.model_fields["profile"].annotation))
tests\test_config_deployment.py:85:     literal_targets =
set(typing.get_args(DeploymentConfig.model_fields["compute_target"].annotation))
tests\test_config_deployment.py:95:     merged = {**AppConfig().model_dump(mode="json")}
tests\test_config_wasb.py:43:     fields = IndexingConfig.model_fields
tests\test_config_wasb.py:57:     assert "ai_chunk_scale_width" in IndexingConfig.model_fields
tests\test_config_wasb.py:64:     fields = IndexingConfig.model_fields
tests\test_reporting.py:55:     def model_dump(self, **kwargs):
tests\test_storage.py:328:     assert isinstance(storage.get_storage(config=sc.model_dump()),
LocalStorage)
tests\test_worker.py:73:     def model_dump(self):
```

14. user (2026-06-21T07:24:16.196Z)

```
Found 21 files
tests\test_config_deployment.py
backend\config\presets.py
backend\config\__init__.py
backend\config\loader.py
backend\config\models.py
docs\archives\decisions\DECISIONS.md
docs\COMPUTE_DECOUPLED_SERVING\README.md
docs\UI_ALPHA_EXECUTION_PLAN.md
docs\PER_VIDEO_WORKSPACE.md
docs\BACKLOG.md
docs\README.md
docs\data_pipeline\GOLDEN_SET_PHASE1_VIDEOS.md
docs\PLATFORM_ARCHITECTURE.md
docs\EXPERIMENT_HARNESS.md
docs\COMPUTE_DECOUPLED_SERVING\architecture.svg
backend\utils\workspace.py
artifacts\appliance_design.json
tests\test_windowing.py
backend\eval\windowing.py
backend\eval\per_user_regression.py
backend\utils\paths.py
```

15. assistant (2026-06-21T07:24:20.381Z)

The `deployment`/`compute\_target` matches in non-config files are likely coincidental (e.g. "deployment" in docs, or unrelated `compute\_target` words). Let me verify the few code-file matches and check `backend/reporting/\_\_init\_\_.py` which dumps config, plus look at any place that snapshots the full config dump or asserts an exact field set.

16. user (2026-06-21T07:24:21.741Z)

```
30         video_id: str,
31         video_path: Optional[str],
32         config: Optional[Dict[str, Any]],
33         val_results: List[Any],
34         val_context: Optional[Dict[str, Any]] = None,
35         segmenter_requested: Optional[str] = None,
36         segmenter_actual: Optional[str] = None,
37         was_reingest: bool = False,
38         segmentation_ran: bool = True,
```

```

39     ) -> Optional[Dict[str, str]]:
40         """Harvest + write the report. Returns the written file paths, or None if
41         reporting is unavailable/disabled/failed. Never raises."""
42         config = config or {}
43         # Accept either a plain dict or a typed config object (e.g.
backend.config.AppConfig).
44         # Normalize to a plain nested dict so harvest's dict access works uniformly.
45         if hasattr(config, "model_dump"):
46             try:
47                 config = config.model_dump(mode="json")
48             except Exception:
49                 config = {}
50         # Consume run signals unconditionally (even when reporting is off) so the
51         # process-local store never accumulates entries for finished runs.
52         signals = run_signals.pop(video_id)
53         if not OBSREPORT_AVAILABLE or not report_enabled(config):
54             return None
55         try:
56             play_segments = db.query_play_segments(video_id, {}) if db is not None else []
57             candidates = db.query_candidate_segments(video_id) if db is not None else []
58             rec = harvest.record_run(
59                 video_id=video_id,
60                 video_path=video_path,
61                 config=config,
62                 val_results=val_results,
63                 val_context=val_context,
64                 play_segments=play_segments,
65                 candidates=candidates,
66                 segmenter_requested=segmenter_requested,
67                 segmenter_actual=segmenter_actual,
68                 was_reingest=was_reingest,
69                 segmentation_ran=segmentation_ran,

```

17. user (2026-06-21T07:24:22.699Z)

```

80         _schedule_next_drain_if_waiting,
81     )
82
83     # Legacy thin wrapper for any remaining direct calls during transition.
84     # New code should import load_app_config / AppConfig from backend.config directly.
85     def load_config(config_path: str = "config.json") -> Dict[str, Any]:
86         cfg = load_app_config(config_path)
87         return cfg.model_dump(mode="json")
88
89     # --- API Models ---
90     class ProcessRequest(BaseModel):
91         video_path: str
92         player_pool: Optional[List[str]] = None
93         max_frames: Optional[int] = None
94         segmenter: Optional[str] = None

```

18. user (2026-06-21T07:24:23.924Z)

```

backend\config\models.py:322:class DeploymentConfig(BaseModel):
backend\config\models.py:328:    ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var);
auto-detect only ever
backend\config\models.py:339:    compute_target: Literal["", "local_gpu", "cloud_run",
"cpu_mock"] = ""
backend\config\models.py:362:    deployment: DeploymentConfig =
Field(default_factory=DeploymentConfig)
backend\config\presets.py:3:A *deployment profile* is a single switch that applies a coherent
BUNDLE of existing
backend\config\presets.py:13:(``config.json`` ``deployment.profile`` or the
``DEPLOYMENT_PROFILE`` env var); with no
backend\config\presets.py:25:# Canonical enum values, defined ONCE here.

```

```

``models.DeploymentConfig``s ``Literal`` mirrors
backend\config\presets.py:26:# these and ``tests/test_config_deployment.py`` pins parity so the
two can never drift.
backend\config\presets.py:30:# Each profile maps 1:1 to its canonical compute_target. Used to
apply the preset and to warn
backend\config\presets.py:31:# when config.json explicitly overrides compute_target into a state
inconsistent with the profile.
backend\config\presets.py:39:# that EXIST today ? M1 is pure/additive: the new
``deployment.{profile,compute_target}`` plus
backend\config\presets.py:45:         "deployment": {"profile": "truly-local", "compute_target":
"local_gpu"},
backend\config\presets.py:52:         "deployment": {"profile": "cloud-serving",
"compute_target": "cloud_run"},
backend\config\presets.py:58:         "deployment": {"profile": "cpu-mock", "compute_target":
"cpu_mock"},
backend\config\presets.py:100:     """SUGGEST (never auto-apply ? D15) a deployment profile from
the environment.
backend\config\__init__.py:13:     DeploymentConfig,
backend\config\__init__.py:26:     "DeploymentConfig",
backend\config\loader.py:48:     """The EXPLICITLY chosen deployment profile, env winning over
config.json. `````` if none
backend\config\loader.py:54:     dep = file_data.get("deployment")
backend\config\loader.py:132:     # consumption sites downstream). A *deployment profile* is opt-
in ? it is applied ONLY when
backend\config\loader.py:133:     # explicitly selected (DEPLOYMENT_PROFILE env or config.json
deployment.profile). With no
backend\config\loader.py:144:         "[CONFIG] unrecognized deployment profile %r (valid:
%s); applying no preset.",
backend\config\loader.py:152:     # so re-stamp it onto the merged deployment section to keep
the record honest.
backend\config\loader.py:153:     merged.setdefault("deployment", {})
backend\config\loader.py:154:     if isinstance(merged["deployment"], dict):
backend\config\loader.py:155:         merged["deployment"]["profile"] = profile
backend\config\loader.py:156:         ct = merged["deployment"].get("compute_target")
backend\config\loader.py:160:         "[CONFIG] deployment profile %r normally uses
compute_target=%r but "
backend\config\loader.py:164:         logger.info("[CONFIG] deployment profile ACTIVE: %s
(compute_target=%s)",
backend\config\loader.py:165:                     profile,
merged["deployment"].get("compute_target"))
backend\config\loader.py:172:     "[CONFIG] no deployment.profile set (manual knobs).
Environment suggests %r ? "
backend\config\loader.py:173:     "set deployment.profile or DEPLOYMENT_PROFILE to
apply (auto-detect never "
backend\utils\workspace.py:9:LOCAL deployment and a CLOUD (bucket) deployment share the *same
relative layout* ? only the root
backend\utils\paths.py:42:     allowed_root (a future hosted/tenant deployment) enforces
containment.
backend\eval\per_user_regression.py:36:``None`` = the local single-user deployment; design-for-
north-star, implement-for-current) and
backend\eval\per_user_regression.py:115:     deployment, keyed as the literal ``"local"``).
``control_spec_hash`` pins WHICH incumbent
backend\eval>windowing.py:18:operating point a given deployment truly serves.
backend\eval>windowing.py:151:     (``"wasb"`` / ``"fusion"`` / ``"tracknet"``). Use this when a
deployment overrode a served

```

19. assistant (2026-06-21T07:24:27.287Z)

Nothing in core code reads `deployment`/`compute\_target` except the config package itself, confirming M1 is inert. Now let me check for any place that snapshots the whole config dump as a fixture/golden, or constructs `AppConfig` with `extra='forbid'`, or asserts an exact top-level field set. Let me also check the worker `--set` precedence claim and whether anything counts top-level sections.

20. user (2026-06-21T07:24:28.873Z)

```

tests\test_ablation.py:274:    forbidden = {"irrelevant", "useless", "dead", "remove"}
tests\test_ablation.py:276:        assert not any(bad in v.lower() for bad in forbidden)
tests\test_base.py:68:        # against "real db/config being passed at discovery time", not
to forbid a
tests\test_config_deployment.py:76:    literal_profiles =
set(typing.get_args(DeploymentConfig.model_fields["profile"].annotation))
tests\test_config_deployment.py:85:    literal_targets =
set(typing.get_args(DeploymentConfig.model_fields["compute_target"].annotation))
docs\RALLY_QUALITY_RESEARCH.md:575:    paid LLMs are inference/serving/shadow-reference only. Also
forbidden: public datasets known to
tests\test_golden_fixtures_split.py:4:the real-fixture-only names (`freeze_real_fixture`,
`RefusalError`, `_scan_forbidden` and the
tests\test_golden_fixtures_split.py:7:every existing import path (`gf.freeze_real_fixture` /
`gf.RefusalError` / `gf._scan_forbidden`)
tests\test_golden_fixtures_split.py:34:def test_scan_forbidden_reexport_is_same_object():
tests\test_golden_fixtures_split.py:35:    """`golden_fixtures._scan_forbidden` is the exact
function from `golden_real_fixtures`."""
tests\test_golden_fixtures_split.py:36:    assert gf._scan_forbidden is grf._scan_forbidden
tests\test_golden_fixtures_split.py:81:def
test_scan_forbidden_behaviour_identical_via_both_paths():
tests\test_golden_fixtures_split.py:82:    """The re-exported `_scan_forbidden` raises the
same `RefusalError` on a forbidden key
tests\test_golden_fixtures_split.py:85:    for fn in (gf._scan_forbidden, grf._scan_forbidden):
tests\test_golden_real_fixtures.py:131:    for forbidden in gf.FORBIDDEN_KEYS:
tests\test_golden_real_fixtures.py:132:        assert forbidden not in exp_text,
f"FORBIDDEN_KEYS leak {forbidden!r} in expected.json"
tests\test_golden_real_fixtures.py:235:def
test_scan_forbidden_catches_label_csv_source_stem(tmp_path):
tests\test_golden_real_fixtures.py:242:    gf._scan_forbidden(text, where="x/expected.json",
tests\test_golden_real_fixtures.py:247:def
test_scan_forbidden_allows_exact_schema_literal(tmp_path):
tests\test_golden_real_fixtures.py:252:    gf._scan_forbidden(text, where="x/expected.json",
source_paths=["/some/dir/trajectory.csv"])
tests\test_golden_real_fixtures.py:458:def test_scan_forbidden_raises_on_forbidden_key():
tests\test_golden_real_fixtures.py:461:    gf._scan_forbidden('{"person": [1,2,3]}',
where="x", source_paths=[])
tests\test_golden_real_fixtures.py:462:    assert "forbidden key" in str(ei.value)
tests\test_golden_real_fixtures.py:465:def test_scan_forbidden_raises_on_md5():
tests\test_golden_real_fixtures.py:469:    gf._scan_forbidden('{"slug": "%s"}' % md5,
where="x", source_paths=[])
tests\test_golden_fixtures.py:57:    for forbidden in gf.FORBIDDEN_KEYS:
tests\test_golden_fixtures.py:58:        assert forbidden not in raw, (
tests\test_golden_fixtures.py:59:            f"{slug}: forbidden key substring {forbidden!r}
present in expected.json (Tier-0 leak)")
docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:27:- **Never train on paid-LLM (Gemini/OpenAI/Anthropic)
outputs.** Also forbidden: public grounding-CoT datasets that
docs\OWNED_MODEL_IMPLEMENTATION_PLAN.md:57:counting; the strong teachers are license-forbidden).
**But build the *seam* now** (the brief requires it):
docs\GOLDEN_REGRESSION_FIXTURES.md:97:- **Strict-slug `label` + dropped `source_note`** ? the
operator `label` is validated against `^[a-z0-9][a-z0-9\-\_]*$` (a free-text identifier like a
player/event name is refused), and operator free-text `source_note` is **accepted-but-NOT-
persisted** (never reaches the committed `provenance.json`) ? both close the free-text
identifier leak `_scan_forbidden` could not police (audit fix #1).
docs\GOLDEN_REGRESSION_FIXTURES.md:101:- **Stream minimization (streams, not resolution)** ?
Tier-0 keeps the shuttle block only; person/audio excluded; pose/gait/face have **no field in
the schema by design** (positive `_scan_forbidden` assertion on the written bytes).
docs\COMPUTE_DECOUPLED_SERVING\TESTING_STRATEGY.md:54:[Omitted long matching line]
backend\eval\golden_real_fixtures.py:12:defined here (`freeze_real_fixture`, `RefusalError`,
`_scan_forbidden`) at the bottom of that
backend\eval\golden_real_fixtures.py:90:def _scan_forbidden(text: str, *, where: str,
source_paths: Sequence[str]) -> None:
backend\eval\golden_real_fixtures.py:103:        f"{where}: forbidden key substring
{key!r} found in written output ? "
backend\eval\golden_real_fixtures.py:185:    back-channel `_scan_forbidden` cannot fully
police (an identifying note that happens to

```

```

backend\eval\golden_real_fixtures.py:213:    # text (a name/event), which _scan_forbidden cannot
reliably catch (red-team #1 free-text leak).
backend\eval\golden_real_fixtures.py:272:    # _scan_forbidden cannot fully police). The fixed
boilerplate `note` constant below stays.
backend\eval\golden_real_fixtures.py:326:        _scan_forbidden((tmp_dir /
"expected.json").read_text(encoding="utf-8")),
backend\eval\golden_real_fixtures.py:328:        _scan_forbidden((tmp_dir /
"provenance.json").read_text(encoding="utf-8")),
backend\eval\golden_real_fixtures.py:332:        _scan_forbidden(str(source_note),
backend\eval\golden_fixtures.py:80:#: here on purpose ? the spec forbids depending on the named
``windowing.SERVED`` constant so a
backend\eval\golden_fixtures.py:115:#: attribution back-channel that ``_scan_forbidden`` does
NOT catch (it isn't a FORBIDDEN_KEY / MD5
backend\eval\golden_fixtures.py:834:    "RefusalError", "freeze_real_fixture",
"_scan_forbidden", "_new_surrogate_slug",
docs\archives\research\OWNED_MODEL_TRAINING_STUDY.md:129:- **Existing distilled model was
trained on Gemini silver labels** ? the (correct) eval-only policy now **forbids
docs\archives\MONETIZATION_AUDIT.md:24:| **MonoTrack** (shuttle) | Adobe Research License ?
**noncommercial only** | N/A ? commercial use forbidden outright | ? **BLOCKER** | Do not use.
Use WASB / TrackNetV3 instead |
docs\archives\MONETIZATION_AUDIT.md:25:| **YOLO-NAS** weights | Deci "YOLO-NAS License" ? bans
commercial/production/managed-service use | N/A ? forbidden | ? **BLOCKER** | Do not use the
released weights |
docs\archives\past_projects\DECOUPLED_COMPUTE\06-RISKS-AND-OPEN-QUESTIONS.md:48:- **DPA / no-
train gate on the *inference-on-uploads* half specifically.** Plans forbid sending non-consented

```

21. user (2026-06-21T07:24:31.445Z)

```

artifacts\so.json:1:[Omitted long matching line]
eval_baselines\experiment_leaderboard.json:103:[Omitted long matching line]
eval_baselines\experiment_leaderboard.json:104:    "served_change": "config.json
indexing.default_segementer gemini->fusion_hybrid; FusionConfig.min_crossings default 0->2",

```

22. assistant (2026-06-21T07:24:42.283Z)

No `extra='forbid'` and no test asserting an exact top-level field set. The `model\_fields` iterations in `models.py` (the `.get()` shim) iterate dynamically, so a new section is handled gracefully ? actually let me verify the `.get()` shim doesn't break: it iterates `type(self).model\_fields` to find a leaf key inside any section. Adding `deployment` adds a section with `profile`/`compute\_target` leaves. Could that shadow an existing leaf key lookup? `profile` and `compute\_target` are new names, no collision. Good.

Now let me verify the key backward-compat claim directly by running the test suites and checking that `profile=''` yields a byte-identical dump. Let me also confirm the worker `--set` precedence and whether the loader's profile path is reachable from the actual repo `config.json`.

23. user (2026-06-21T07:24:52.204Z)

```

..... [100%]
52 passed in 6.22s

```

24. assistant (2026-06-21T07:25:00.535Z)

Both suites pass. Now let me verify the precise byte-identity claim and probe the precedence/edge cases that the tests may not cover. Specifically:

1. The repo `config.json` ? does it have a `deployment` section? If not, does `load\_config` on it produce a config where `deployment` is well-formed (the `else` branch)?
2. The precedence claim "config.json < env" ? for the *non-profile* knobs, env overrides happen "downstream" not in the loader. Is `DEPLOYMENT\_PROFILE` the only env read in the loader? Yes.
3. Critically: does the profile branch's `\_deep\_merge(defaults, file\_data)` change behavior vs the no-profile shallow merge for *sections* the user provided? The no-profile path is shallow (replaces whole section); the profile path is deep-merge. This is intended and tested. But is

there a subtle bug: when a profile is active, a user's `config.json` `indexing` section now deep-merges instead of shallow-replacing. That's the documented behavior change *only* when a profile is opted into*, so it doesn't affect default behavior.

Let me check the actual repo config.json and run the full broader test sweep to confirm nothing else breaks.

25. user (2026-06-21T07:25:02.575Z)

```
top-level keys: ['ai_model', 'ai_provider', 'indexing', 'output', 'sport', 'stitching',
'validation']
has deployment: False
default_segmenter: gemini
```

26. assistant (2026-06-21T07:25:05.731Z)

The repo `config.json` has no `deployment` section and no profile ? so the production path takes the `else` (no-profile) branch and is byte-identical to pre-M1. Let me verify the byte-identity of the produced config against a simulated pre-M1 state, and confirm importing `presets` doesn't import torch or do IO.

27. user (2026-06-21T07:25:11.694Z)

```
torch imported at module load: False
profile: '' compute_target: ''
default_segmenter (must stay gemini): gemini
detector_impl (must stay auto): auto
storage.backend (must stay local): local
deployment section value: {'profile': '', 'compute_target': ''}
```

28. assistant (2026-06-21T07:25:17.769Z)

No torch import at config load, and the no-profile path is unchanged. Now let me probe two precedence edge cases that could be subtle bugs, since the task asks me to ground claims with exact triggering inputs:

1. **The unknown-env-profile + valid-config-profile interaction.** If `DEPLOYMENT\_PROFILE=banana` (env, invalid) AND `config.json` has `deployment.profile=cpu-mock` (valid), what happens? `\_resolve\_explicit\_profile` returns env "banana" (env wins). Then `is\_known\_profile("banana")` is False ? warns, sets `profile=""` ? takes the **else** (no-profile) branch. So the valid config.json profile is silently ignored because the env had a typo. Let me verify this and judge whether it's a real issue.
2. **The `compute\_target` re-stamp / mismatch warning** when the env selects a different profile than config.json which set a `compute\_target`.

29. user (2026-06-21T07:25:21.721Z)

```
[CONFIG] unrecognized deployment profile 'banana' (valid: truly-local, cloud-serving, cpu-mock);
applying no preset.
CASE1 invalid-env + valid-file-profile:
  profile: 'cpu-mock' (expected cpu-mock if file honored, "" if env-typo clobbers)
  detector_impl: auto (stub=preset applied, auto=no preset)
  ai_max_workers: 9
```

30. assistant (2026-06-21T07:25:31.177Z)

This confirms a real edge case. When `DEPLOYMENT\_PROFILE` has a typo (`banana`) **and** `config.json` validly sets `deployment.profile=cpu-mock`:

- The loader warns about the bad env value and sets `profile=""`, taking the **no-profile else** branch ? **no preset is applied** (`detector\_impl` stays `auto`, not `stub`).
- But then `AppConfig.model\_validate(merged)` runs on the shallow-merged dict which still contains `file_data["deployment"]["profile"] = "cpu-mock" ? so the *recorded*

`cfg.deployment.profile` is `"cpu-mock"`, while the **behavior** (detector_impl/skip_ai_handoff/storage) is ***default/unconfigured***.

The result is an ***incoherent config***: `cfg.deployment.profile == "cpu-mock"` but none of the cpu-mock preset knobs are applied. A user reading `cfg.deployment.profile` would believe the profile is active when it is not. This is a real correctness issue, though it requires a fairly specific misconfiguration (env typo + valid file profile). It is not a backward-compat / "zero core change" violation (the default `profile=''` path is unaffected), so severity is minor.

Let me double check by confirming what the docstring/intent says ? line 138-147 says "A bad value coming only from the env override is ignored here with this warning." The intent is to **ignore** the bad env value, but the implementation discards the **fallback to config.json** entirely instead of falling back to it. Let me verify the fallback expectation isn't covered elsewhere and check whether `\_resolve\_explicit\_profile` could fall back.

31. user (2026-06-21T07:25:32.197Z)

```
47     def _resolve_explicit_profile(file_data: dict[str, Any]) -> str:
48         """The EXPLICITLY chosen deployment profile, env winning over config.json. ```` if
49         none
48         is set. Auto-detect is deliberately NOT consulted here ? it only suggests (D15), so
49         a bare
50         environment never silently selects a profile (and never silently spends on
51         cloud)."""
51         env_profile = os.environ.get("DEPLOYMENT_PROFILE")
52         if env_profile and env_profile.strip():
53             return env_profile.strip()
54         dep = file_data.get("deployment")
55         if isinstance(dep, dict):
56             prof = dep.get("profile")
57             if isinstance(prof, str) and prof.strip():
58                 return prof.strip()
59         return ""
```

32. assistant (2026-06-21T07:25:39.295Z)

Confirmed. `\_resolve\_explicit\_profile` returns the env value as soon as it's non-blank, without validating it. A bad env value then short-circuits and the valid config.json profile is never consulted. The "ignored here" comment is aspirational ? the env value isn't really ignored, it actively suppresses the file's valid profile. This is the minor finding.

Now let me check one more subtle thing for the backward-compat lens: the `else` branch calls `suggest\_profile()` unconditionally on every default load. This calls `os.environ` reads (`K\_SERVICE`, `CI`) ? harmless, no torch (probe_cuda defaults False). Confirmed cheap. But does it have a side effect that changes behavior when the var is unset? No ? it just logs at DEBUG level. Confirmed inert.

Let me also verify the `validate\_assignment=True` interaction: the loader does `merged["deployment"]["profile"] = profile` on a plain dict before `model\_validate`, so no assignment validation triggers there. Fine.

Let me run the broader test suite to make sure no snapshot/golden test breaks from the new section, and check the worker tests since they `model\_dump()` storage config.

33. user (2026-06-21T07:25:48.910Z)

```
..... [ 90%]
..... [100%]
80 passed in 6.58s
```

34. assistant (2026-06-21T07:25:56.090Z)

All green. Let me do a final check on the precedence claim in the spec vs implementation. The spec §4.3 line 84 says precedence is "built-in defaults ? profile preset ? config.json ? env ?

worker --set". The loader implements defaults < preset < config.json for the *bundle knobs*. But note: the env (`RALLY\_DETECTOR\_IMPL`, `WASB\_\*`) overrides happen downstream at consumption sites, not in the loader ? so for those, env > config.json still holds outside the loader. `DEPLOYMENT\_PROFILE` is the one env handled in-loader, and it only affects *which profile* is chosen, not individual knobs. This is consistent.

One last verification: does the profile-active path correctly let env knob overrides still win? E.g., profile=cpu-mock sets detector_impl=stub, but `RALLY\_DETECTOR\_IMPL=native` env should still override. That's a consumption-site concern (not in this loader), and M1 doesn't touch those sites. Let me confirm the detector_impl env override is read downstream and not regressed.

35. user (2026-06-21T07:25:57.103Z)

```
backend\config\models.py:132:    # RALLY_DETECTOR_IMPL env var overrides this. See
pipeline\detectors\{stub,native_wasb}_runner.py.
backend\worker\__main__.py:9:are byte-for-byte unchanged. Combine with
`RALLY_DETECTOR_IMPL=stub` and
backend\worker\__main__.py:141:    detector_impl =
(os.environ.get("RALLY_DETECTOR_IMPL")
backend\worker\__main__.py:214:    detector_impl (RALLY_DETECTOR_IMPL / indexing.detector_impl:
auto=WSL, native=GPU box,
backend\pipeline\detectors\stub_runner.py:20:`indexing.detector_impl="stub"` or the
`RALLY_DETECTOR_IMPL=stub` env var; the real
backend\pipeline\segmenters\trajectory_hybrid.py:68:    `RALLY_DETECTOR_IMPL` env var
(takes precedence) or `indexing.detector_impl`:
backend\pipeline\segmenters\trajectory_hybrid.py:76:    impl =
(os.environ.get("RALLY_DETECTOR_IMPL") or idx_cfg.get("detector_impl") or "auto").lower()
backend\pipeline\segmenters\trajectory_hybrid.py:190:    # _resolve_runner honours the CPU-
stub override (RALLY_DETECTOR_IMPL /
```

36. assistant (2026-06-21T07:26:08.384Z)

Confirmed: `RALLY\_DETECTOR\_IMPL` env override is read at the consumption sites (`trajectory\_hybrid.py:76`, `worker/\_\_main\_\_.py:141`) and is untouched by M1, so the env > config.json > preset precedence holds end-to-end. M1 is purely additive at the config-loading layer.

I have completed a thorough review. Let me summarize my assessment.

```
**Backward-compat / "zero core change" verdict (PASS):**
- `profile=''` (default) takes the `else` branch (loader.py:166-175): the original shallow
`{**defaults, **file_data}` merge is preserved byte-for-byte. Verified empirically ? the repo
`config.json` has no `deployment` section/profile, loads via the else branch, and
`default_segmenter`/`detector_impl`/`storage.backend` are unchanged.
- `test_default_config_dump_only_adds_deployment_section` (test_config_deployment.py:62) plus
`test_non_object_config_degrades_to_defaults` and `test_config_shallow_section_replacement`
(test_config.py) all pass ? the prior contract holds.
- No caller uses `extra='forbid'` or asserts an exact top-level field set; nothing in core reads
`deployment`/`compute_target`. The `AppConfig.get()` shim (models.py:376-394) iterates
`model_fields` dynamically, so the new section is handled gracefully and the new leaf names
(`profile`, `compute_target`) don't shadow any existing leaf.
- `import backend.config.presets` imports no torch and does no IO (torch is lazy in
`_cuda_available`, only reached via opt-in `probe_cuda=True`). Verified `torch` not in
sys.modules` after load.
- `load_config` reading `os.environ['DEPLOYMENT_PROFILE']` when unset:
`_resolve_explicit_profile` returns `""`, no behavior change. `suggest_profile()` on the default
path is env-only/cheap and only logs at DEBUG.
```

****One minor finding**** (not a backward-compat violation): a typo'd `DEPLOYMENT\_PROFILE` env value suppresses a valid `config.json` profile, yielding an incoherent config where `cfg.deployment.profile` is recorded as the file's valid value but no preset is applied.

37. user (2026-06-21T07:26:21.142Z)

Structured output provided successfully

38. assistant (2026-06-21T07:26:31.201Z)

I've completed the M1 review. Summary:

****Backward-compat / "zero core change" lens: PASS.**** The hard requirement holds ? `profile=''` yields a config byte-identical to pre-M1:

- The no-profile path takes the original shallow `{\*\*defaults, \*\*file\_data}` merge (loader.py:166-168), verified empirically against the real repo `config.json` (no `deployment` section ? else branch ? `default\_segmenter`/`detector\_impl`/`storage.backend` all unchanged).
- Adding the `deployment` section breaks no caller: no `extra='forbid'`, no test asserting an exact top-level field set, nothing in core reads `deployment`/`compute\_target`, and the `AppConfig.get()` shim iterates `model\_fields` dynamically (new leaf names don't collide).
- `import backend.config.presets` imports no torch and does no IO (torch is lazy, gated behind opt-in `probe\_cuda=True`).
- Unset `DEPLOYMENT\_PROFILE` causes no behavior change; `suggest\_profile()` on the default path is env-only and only logs at DEBUG.
- The prior contracts `test\_config\_shallow\_section\_replacement` and `test\_non\_object\_config\_degrades\_to\_defaults` still pass; all 52 deployment + 80 config-adjacent tests green.

****One minor finding**** (high confidence): a typo'd `DEPLOYMENT\_PROFILE` env value suppresses an otherwise-valid `config.json` profile, producing an incoherent config (`cfg.deployment.profile` recorded as the valid file value while the preset is not actually applied). This is a robustness gap in the new opt-in precedence logic, not a backward-compat violation. Details and a suggested fix are in the structured output.

Relevant files: `C:/Users/avidu/Projects/badminton-highlight-indexer/backend/config/loader.py`,
`C:/Users/avidu/Projects/badminton-highlight-indexer/backend/config/presets.py`,
`C:/Users/avidu/Projects/badminton-highlight-indexer/backend/config/models.py`.