

claude-windows / dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3\subagents\workflows\wf\_f1431155-4a0\agent-adcdea673ac23f4c0.jsonl`
- SHA-256: `c0f704e346da843e6bcce8a1314dff4c128b7c1ffd4bed61f8f4b9b62cffa801`
- Source modified: `2026-06-21T07:26:22+00:00`
- Imported at: `2026-07-05T16:48:22+00:00`
- project: `wf\_f1431155-4a0`
- session_id: `dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3`

Transcript

1. user (2026-06-21T07:23:57.138Z)

You are reviewing milestone M1 of the COMPUTE_DECOUPLED_SERVING project in the repo at C:/Users/avidu/Projects/badminton-highlight-indexer.
The spec is docs/COMPUTE_DECOUPLED_SERVING/README.md (read §2 D7-D15, §4.3 "Profile selection", and §7 the tracker row M1). M1 must be: PURE, ADDITIVE, with deployment.profile='' (the default) == today's behaviour, ZERO core change. A *deployment profile* is opt-in and applies a coherent preset bundle of EXISTING knobs (indexing.detector_impl/skip_ai_handoff, storage.backend) plus the new deployment.compute_target. Precedence: built-in defaults < profile preset < config.json < env.
D15: auto-detect SUGGESTS, never auto-applies (cloud spend never entered silently).

Files changed (read them all, plus the existing tests/test_config.py to understand the prior contract):

- backend/config/presets.py (NEW)
- backend/config/models.py (added DeploymentConfig + AppConfig.deployment field)
- backend/config/loader.py (added _deep_merge, _resolve_explicit_profile, preset application)
- backend/config/__init__.py (re-exports)
- tests/test_config_deployment.py (NEW)

Ground EVERY claim in the actual code with a file:line citation. Do NOT speculate. If you assert a bug, give the exact input that triggers it and the wrong-vs-right output. Rate confidence honestly; prefer FEWER, HIGH-CONFIDENCE findings. A clean bill ([] findings) is a valid and good result.

LENS: test quality & coverage gaps in tests/test_config_deployment.py. For each behaviour in the loader, is there a test that would FAIL if that behaviour regressed? Hunt for tests that pass for the wrong reason (e.g. asserting a value that's the default anyway, so it wouldn't catch a missing preset).
Check the torch-free guarantee is actually tested, the env+file conflict is tested, the unknown-profile (file=raise vs env=warn) split is tested, and the deep-merge-sibling-survives case is tested.
Name any missing high-value test case with the exact assertion it should make.

2. assistant (2026-06-21T07:23:59.588Z)

I'll review M1 of the COMPUTE_DECOUPLED_SERVING project, focusing on test quality and coverage gaps. Let me start by reading the spec and the changed files.

3. user (2026-06-21T07:24:01.414Z)

```

1      # Compute-Decoupled Rally Serving ? one pure core, profile-selected compute + storage
2
3      > **Status:** ? **M0 APPROVED** (owner, 2026-06-21) ? **IN PROGRESS (build M1+)** .
**Owner:** Avi . **Created:** 2026-06-21 . **Last updated:** 2026-06-21
4      > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to
`docs/archives/past_projects/` when DONE)
5      > **Tracking anchors:** $7 progress tracker is the source of truth; pointer in memory
`current-state` + `docs/NEXT_STEPS.md` once work starts.
6      > **Relation to existing docs:** the productionization successor to the archived
`DECOUPLED_COMPUTE` (M3/M4); realizes the `ComputeBackend`/`StorageBackend` registries from
[PLATFORM_ARCHITECTURE.md](../PLATFORM_ARCHITECTURE.md) $3; the **truly-local** profile is the
L0 tier of [VIDEO_LOCALITY_MODEL.md](../VIDEO_LOCALITY_MODEL.md) and the `LocalDesktop` backend
of [DESKTOP_APP_PLAN.md](../DESKTOP_APP_PLAN.md).
7      > **Naming:** this was the "Cloud Run + GPU serving subproject" (ADR-012). Renamed
**compute-decoupled serving** because **truly-local is a co-equal, first-class profile**, not an
afterthought.
8      > **Honesty note:** claims tagged `[verified]` (checked against code by the design
workflow `wf_56274ff0`), `[design]` (proposed), `[researched]`.
9
10     ---
11
12     > **? Northstar (D14):** a **mobile, app-like** flow ? point at a **Google Drive** video
? **fully cloud-native** run ? the stitched, **ready-to-share** reel lands **back in Drive, next
to the source**. Two design promises around it: a user can **switch local?cloud anytime** (D13,
a headline feature), and execution is **never a surprise** ? auto-detect **suggests**, the user
**confirms** (D15).
13
14     ## 0. TL;DR
15     The rally engine is a **pure function of `(input bytes + config) ? outputs`** in three
deterministic stages ? **DETECT ? WINDOW ? STITCH** ? that **no deployment profile may branch
on**. A **`ComputeBackend` (deployment profile)** decides **where** the core runs (a local process
vs a Cloud Run GPU job); a **`StorageBackend`** decides **where bytes come from** (local FS / USB
/ mounted Google Drive / bucket). **One startup config** picks both. We ship **truly-local**
(local GPU + in-process stitch, zero cloud) and **cloud-serving** (upload `<2GB` / gdrive /
bucket ? a Cloud Run GPU job runs detect **and** stitch, so stitch's compute is metered), plus
**cpu-mock** for CI. The single most important caveat: detection is **byte-identical only on the
same GPU** (cross-GPU is sub-pixel-different, per the 2026-06-20 deploy report) ? so the
**parity guarantee is enforced at the WINDOW + STITCH layer**, not detector-CSV bytes.
16
17     ## 1. Problem & goal
18     - **Why now:** DECOUPLED_COMPUTE proved compute **portability** (M0?M2 + M3.5 on a GCP L4)
and is archived; its deferred **M3 (serving endpoint + per-video billing) + M4 (scale-to-zero +
cost guard)** become this project. The owner needs **both** a hosted service **and** a truly-local
mode, with the **core (detect+stitch) unaffected** by which one runs.
19     - **The two user-facing outcomes** (owner, 2026-06-21):
20     - ** (a) Hosted service** ? operate on small **local uploads (`<2GB`)** **and** load from
**gdrive / a bucket**, spinning up **Cloud Run for both rally segmentation AND stitching**
(stitching is compute-intensive ? run it where its cost is accounted for).
21     - ** (b) Truly-local** ? on a local machine with videos on **local drives/USB + a local
GPU**, run inference **and** stitching **completely locally**.
22     - **"Good" looks like:** the same `video ? rallies ? reel` core gives identical windows
+ stitched ranges whether it runs on a laptop or Cloud Run; switching is a **config/startup
choice**, never a code branch in the core.
23
24     ## 2. Decisions locked
25
26     | # | Decision | Source | Implication |
27     |---|-----|-----|-----|
28     | D1 | The core is a **pure 3-stage function** (DETECT/WINDOW/STITCH); **no profile
branch inside it**. | design `[verified]` | All profile difference lives in config + the two
registries. |
29     | D2 | **Two registries:** `ComputeBackend` (where) selected by `deployment.profile` +
`compute_target`; `StorageBackend` (what-bytes) by `input_backend`/`storage.backend`. | ADR-014
| Realizes PLATFORM_ARCHITECTURE's `local`|`hosted`|`byo_worker`. |
30     | D3 | **Profiles shipped:** `truly-local`, `cloud-serving`, `cpu-mock` (CI);

```

`byo\_worker` reserved/deferred. | ADR-014 | Enum/seam reserved so BYO isn't precluded. |
31 | D4 | In **cloud-serving**, STITCH runs on Cloud Run** (co-located with detect) so CPU/wall + egress are metered into the per-job cost ledger. | owner (a) | Stitch-on-laptop would hide real cost. |
32 | D5 | **Parity enforced at WINDOW+STITCH** (byte-exact for a fixed trajectory); **detect** is byte-exact only same-GPU. | deploy report 2026-06-20 | Cross-GPU asserted at the rally/window level, never CSV bytes. |
33 | D6 | **Default-OFF, smallest-safe-first, one PR per milestone**; any core-caller change (e.g. configurable output base) ships behind a flag + a Launch Report. | **[[launch-report-convention]]** | Zero-regression discipline. |
34 | D7 | **Cloud Run cold-start: SCALE-TO-ZERO** ? accept the ~1?3 min cold-start (it's an async job: upload?poll?reel, amortized over the multi-minute job). **No warm pool** (that ~\$500/mo idle would break per-video billing). | owner 2026-06-21 | M6/M7. Revisit a warm lane only on a UX complaint. |
35 | D8 | **Pricebook = a versioned DB table**; cost computed in **USD** (matches GCP billing = one source of truth), **displayed in INR** at a configurable FX rate (Bangalore market); re-versioned when GCP prices change. | owner 2026-06-21 | M8. |
36 | D9 | **Edge auth = Cloudflare Access** (reuse the site's existing CF). Lock the Cloud Run **ingress** to the CF proxy + **verify the `Cf-Access` JWT signature** (so a direct hit to the Cloud Run URL can't spoof the identity header). | owner 2026-06-21 | M9. One identity system. |
37 | D10 | **Free-tier "data-for-reels" trade**: free reels in exchange for **training rights** (uploaded footage + derived trajectories/labels); **paid tier = no-train** (privacy is the paid feature). **Carve-outs**: train ONLY on attested **ADULT** footage (minor footage ? reel yes, training pool **NO** ? DPDP/GDPR need verifiable parental consent); train on **DERIVED** data + **auto-delete raw**; **ToS counsel-reviewed**; live training-on-uploads **gated to M9** (public signup). Truly-local needs no DPA. | owner 2026-06-21 | M9 + D12; ties **[[paper-coauthor-aim]]** / PERSONALIZATION flywheel. |
38 | D11 | **Quality floor: Gemini handoff ON** for cloud-serving until the owned model clears the #172 ship-gate (e.g. ?0.70 multi-court) ? then flip via a **Launch Report**; the owned model runs in **shadow/eval** meanwhile. (Gemini = serving/inference ONLY, never a training source ? CLAUDE.md guardrail.) | owner 2026-06-21 | M7 + D12. |
39 | D12 | **Data-flywheel harness is IN SCOPE** (owner add, Q5): capture the (consented, adult) uploaded video + the Gemini reel/rally output ? **reliably store** in the per-video workspace/bucket ? **run shadow evals** (owned-vs-Gemini, dual-eval-set) ? **promote good ones** to the golden TRAINING set (human-reviewed labels) so the owned model climbs toward the D11 floor (then Gemini flips off). | owner 2026-06-21 | new **M11**; ties D10 (consent) + **data_pipeline/** + the eval/experiment harness. |
40 | D13 | **Profile is user-switchable ANYTIME** ? a headline FEATURE. The same user runs **local today, cloud tomorrow, back to local** ? it's just a config/startup choice; the pure core gives **equivalent** output either way. **Advertise it** ("your machine *or* our cloud ? your call, switch anytime"). | owner 2026-06-21 | Honest caveat: equivalent at the **rally/window** level, not bit-identical across GPUs (D5); a single job runs in one profile, switching is **between** jobs. |
41 | D14 | **NORTHSTAR** ? operate toward this: a **mobile, app-like** flow ? point at a **Google Drive** video ? **fully cloud-native** run ? the stitched, **ready-to-share** reel lands **back in Drive, next to the source**. | owner 2026-06-21 | Implies a **gdrive-api** (OAuth) StorageBackend (read source + write reel back) ? distinct from the local **mount** backend; **resolves \$8 #7**, new **M12**. The mobile UI is the khelsutra track; the engine must support it (async jobs + status + Drive round-trip / signed URL). Net-new scope (OAuth + Drive write); not necessarily v1, but the design must not preclude it (the **StorageBackend** ABC accommodates it). |
42 | D15 | **Auto-detect SUGGESTS**, the user **CONFIRMS** ? execution is NEVER a surprise. A GPU owner may still choose cloud; **cloud (= spend)** is never entered silently; the active profile is always explicit + surfaced. | owner 2026-06-21 | Amends \$4.3 (auto-detect = a proposed default, not an auto-decision). |
43
44 **## 3. Foundation ? evolution / ADR lineage** ? **trace the idea end-to-end**
45 This project is the latest step in a long decision chain; each ADR contributed a piece and a **subtlety that carries forward**:
46
47 | ADR / doc | Contributed | Subtlety carried into this project |
48 |---|---|---|
49 | **ADR-005** | Permissive licensing (MIT/Apache-2.0/BSD only) | The engine + any served model + the **container image** (ffmpeg, fonts, weights) stay commercial-clean. |

```

50      | **ADR-008** | Owned-model topology; **train==serve featurizer single-sourcing** | The
"core unaffected across profiles" is the same parity discipline, extended from train/serve to
**local/cloud**. |
51      | **ADR-009** | KhelSutra Guru; **local-first delivery** | The **truly-local profile**
is the direct continuation ? privacy/offline self-host stays first-class. |
52      | **ADR-010** | DECOUPLED: D0?GCP pivot | The cloud-compute origin (GPU + co-located
bucket). |
53      | **ADR-011** | DECOUPLED scope = M0?M2+M3.5; **deferred M3 (serving) + M4
(billing/cost-guard)**; overrode the *"rent nothing / not a service"* posture | **This project
realizes M3/M4.** Carries $06's owner-gated gates: **DPA/consent for non-owner uploads,
collector `md5` keying, WASB-C3**. |
54      | **ADR-012** | GPU-native; TPU deferred; **Cloud Run+GPU scale-to-zero = serving
model**; **NVDEC** named as the decode lever | The serving substrate + the 7-step seed scope
this expands. |
55      | **ADR-013** | Drop D0; **GCP sole compute stack**; $200 D0 credits ? non-compute only
| The provider for cloud-serving; D0 Spaces remains only a possible S3 *storage* target. |
56      | **? ADR-014 (new)** | **Compute-decoupled serving:** one pure core, profile-selected
compute+storage; stitch-where-metered | This doc. **Refines** ADR-011/012, does **not**
supersede them. |
57      | PLATFORM_ARCHITECTURE.md | The `ComputeBackend`/`StorageBackend` registries | The
abstraction this realizes in code. |
58      | VIDEO_LOCALITY_MODEL.md / DESKTOP_APP_PLAN.md | L0 fully-local tier / `LocalDesktop`
backend | The truly-local profile = these, productionized. |
59      | 07-COMPUTE-ABSTRACTION (archived) | `DeviceContext`/`DeviceProfile` (designed) | WASB
hardcodes `device='cuda'` with **no CPU fallback** ? a portability gap M4 closes. |
60      | DEPLOY_REPORT_GCP_2026-06-20 (archived) | First real L4 run; **cross-GPU sub-pixel
parity finding** | Why D5 (parity at window level, not CSV bytes). |
61
62      ## 4. Design / architecture
63      See [architecture.svg](architecture.svg) for the picture. **The core never changes;
the box around it does.**
64
65      ### 4.1 The core contract (must stay byte/behaviour-identical) `[verified]`
66      1. **DETECT (GPU):** `DetectorRunner.run_predict(video_path, output_dir, video_id) ?
CSV[Frame,Visibility,X,Y]` (`backend/pipeline/detectors/base.py:164`). Three interchangeable
impls ? `WasbRunner` (WSL), `NativeWasbRunner` (Linux GPU), `StubDetectorRunner` (CPU mock) ?
emit the **identical CSV schema**. The trajectory CSV is the stable artifact boundary.
67      2. **WINDOW (CPU pure fn):** `trajectory_to_action_windows(points, fps, frame_width, ?)
? windows` (`tracknet_runner.py:281`). A shared `@staticmethod`, zero GPU/IO/randomness ? same
trajectory + config = byte-identical candidate windows. Used by every runner *and* the
eval/regression stack verbatim.
68      3. **STITCH (CPU/ffmpeg):** `VideoCompiler.compile_highlights(raw_video, segments,
out_name) ? reel` (`stitcher/compiler.py:303`). Pure FS+subprocess (merge ? ffmpeg ? watermark
? per-clip libx264 trim ? concat). Deterministic for a given input + ranges + watermark config.
69
70      **Non-goal (the trap to avoid):** *no* code branch inside detect/window/stitch keyed on
profile ? the same discipline the experiment harness already enforces (a disabled cue is byte-
identical to CONTROL).
71
72      ### 4.2 The deployment profiles
73      | Profile | Compute | Storage | Orchestration | When |
74      |---|---|---|---|---|
75      | **truly-local** | `compute_target=local_gpu`; `detector_impl=native` (Linux) / `auto`
(WSL); stitch **in-process** | `local` (or `gdrive` = mounted Drive) | `python -m backend.main`
or `worker infer` CLI; existing async job worker | Self-host / offline / privacy / dev-on-a-GPU-
box ? owner (b) |
76      | **cloud-serving** | `compute_target=cloud_run`; `detector_impl=native` (L4, cuda);
**detect AND stitch on the same Cloud Run job** | `gcs` (or gdrive mount / assembled upload);
per-video workspace; reels via signed URL | Cloud Run control plane enqueues ? Cloud Run GPU job
pulls/runs/pushes; `WAIT_AND_RESUME` reschedules the control-plane job; scale-to-zero | Hosted
SaaS ? owner (a) |
77      | **cpu-mock** | `compute_target=cpu_mock`; `detector_impl=stub` (synth/replay) |
`local`/in-memory fakes | `run_all_tests.py` / pytest | CI, regression, profile-parity, GPU-free
dev |
78      | *byo_worker* (deferred) | tenant GPU via outbound pull agent | tenant's own bucket

```

```

(signed URLs) | long-poll pull; same job table scoped by tenant | Privacy-sensitive enterprise ?
**DEFERRED**, enum/seam reserved |
79
80     ### 4.3 Profile selection (the "startup/setup config") `[design]`
81     ONE new top-level `deployment` section on `AppConfig` + preset application + env auto-
detect in `load_config`:
82     - `deployment.profile` ? {truly-local | cloud-serving | cpu-mock} (empty = today's
manual knobs, **fully backward-compatible**).
83     - Selecting a profile applies a **coherent preset bundle** of *existing* knobs that
today drift independently: **INPUT** (new `input_backend`/`input_root`, parallel to the output-
side `StorageConfig`), **COMPUTE** (new `compute_target` enum locking `detector_impl` +
`skip_ai_handoff` + workspace strategy), **STORAGE**
(`storage.backend`/`workspace_root`/`single_folder_per_video`, all already exist).
84     - **Precedence:** built-in defaults ? profile preset ? `config.json` ? env
(`RALLY_DETECTOR_IMPL`, `WASB_*`, `DEPLOYMENT_PROFILE`, ?) ? worker `--set k=v`. Every knob
stays individually overridable.
85     - **Env auto-detect SUGGESTS, the user CONFIRMS (D15 ? no surprise execution):** the env
gives a *proposed* default (`K_SERVICE` ? cloud-serving; `CI=true` ? cpu-mock; else
`torch.cuda.is_available()` ? truly-local), but the user **explicitly confirms or overrides** it
? a GPU owner may still pick cloud, and **cloud (= spend) is NEVER entered silently**. The
active profile is always surfaced (logged/shown). **Per-profile startup checks fail/warn fast**
(cloud-serving without GCS creds/GPU ? loud error before any job).
86     - The worker (`cmd_infer`/`cmd_train`) is **already profile-agnostic** (reads
`config.storage`, accepts `--config`/`--set`) ? no worker rewrite.
87
88     ### 4.4 Compute placement ? and *why stitch runs on Cloud Run* `[design]`
89     DETECT on GPU everywhere; WINDOW CPU-pure everywhere; STITCH is real CPU work. In
**cloud-serving, stitch is co-located in the Cloud Run job** so `gpu_active_seconds` (detect) +
`wall_seconds` (detect+stitch) + `bytes_out` (reel egress) land on **one metered invocation** ?
the per-job cost ledger. Running stitch on a laptop would leave that cost
**invisible/unattributed** ? exactly what the owner wants to avoid. (Stitch stays synchronous in
the GPU job for the common one-reel case; a queued CPU-only stitch fallback only if compiles get
bursty ? the job table + `claim_due_job` already supports both.)
90
91     ### 4.5 Reuse map ? most of the seams already exist `[verified]`
92     **Exists:** the 3-stage core; `_resolve_runner` (the single compute seam);
`StubDetectorRunner` (`replay_csv` = byte-exact $0 anchor); `StorageBackend` ABC + 4 backends +
`StorageRef.parse_uri`; the worker CLI (fetch/put, `--config`/`--set`); the async job control
plane (`jobs` table, `claim_due_job` CAS); `ExecutionPolicy WAIT_AND_RESUME`/`JobWaiting`
(carries to Cloud Run dispatch verbatim); the chunked-upload router (<2GB); `skip_ai_handoff`
(LocalNoAI); storage fakes; golden fixtures + experiment harness. **Partial:** per-video
workspace helper (default-OFF); `DeviceContext` (designed); edge-auth (v4 `user_id` columns
exist). **Net-new:** the `deployment` section; `input_backend`; configurable output base (kill
hardcoded `output/`); Cloud Run dispatch shim + container image; cost ledger + pricebook; edge-
auth header extraction; startup env detection.
93
94     ### 4.6 The data-flywheel harness (D12 / M11) `[design]`
95     Gemini-on (D11) ships good reels now; this harness is **how we earn the right to flip
Gemini off**. On a **consented, adult** cloud-serving job (the D10 trade), the pipeline
additionally:
96     1. **Captures** the source upload + the Gemini reel + the rally output
(intervals/report) into the per-video workspace (`workspaces/<video_id>/`), instead of the no-
consent auto-delete path.
97     2. **Reliably stores** them durably in the bucket (the consented-retention bucket; raw
video kept only for the training pool, not the no-consent path).
98     3. **Runs shadow evals** ? the owned model scores the same video alongside Gemini
(`backend/eval/experiment.compare` + the dual-eval-set, [[eval-dual-set-regression]]) ? tracks
the owned-vs-Gemini gap per job (the live read on "is the floor reachable yet?").
99     4. **Promotes to golden** ? captured videos that pass review become **golden TRAINING
rows** (human-reviewed labels via the annotation flow ? `data_pipeline/`), growing the corpus ?
the owned model climbs to the D11 floor ? flip Gemini off via a Launch Report.
100     This is the consent?capture?eval?goldenize **flywheel** (the cloud realization of the
PERSONALIZATION contribution model): every consented free reel makes the owned model better.
Reuses `data_pipeline/`, `backend/eval/`, and the M9 consent gate; **adults-only**, raw-
retention gated by the D10 consent.

```

```

101
102     ## 6. Honest limits ? what this does NOT do / prove
103     - A green CI suite proves **plumbing + determinism (window+stitch) + profile-parity**
for $0 ? it does **NOT** prove field quality, nor that real-GPU detection is good (that's owner-
side, real-video).
104     - **Cross-GPU detection is NOT byte-identical** ? parity is a window-level claim, not a
CSV-byte claim.
105     - The cloud `0.1.0-l4` weights are **n=2 plumbing-grade** ? cloud-serving ships with
**Gemini handoff ON** until the owned model clears a floor (open Q).
106     - This does not resolve **WASB-C3** (sharing a trained model) or the **DPA/consent**
gate ? both remain owner/counsel-gated before public, non-owner serving.
107
108     ## 7. Deliverables & progress tracker ? **source of truth**
109     Legend: ? Todo · ? In progress · ? Done · ? Gated. **One small PR per row**, off
`origin/master`, no stacking. Default-OFF where it touches core callers.
110
111     | ID | Deliverable | Depends | Gated? | Status | PR |
112     |----|-----|-----|-----|-----|----|
113     | M0 | This design doc + project dir + `runlogs/` + ADR-014 | ? | owner sign-off | ? |
**? APPROVED 2026-06-21** |
114     | M1 | `deployment.profile` + `compute_target` config + preset bundles +
precedence/auto-detect; unit tests (no behaviour change, profile='' default) | M0 | safe | ? | ? |
|
115     | M2 | `input_backend`/`input_root`; wire main CLI/API input through
`StorageRef.parse_uri` (local=identity); storage-fake tests | M1 | safe | ? | ? | ? |
116     | M3 | Configurable output/scratch base ? kill hardcoded `output/`
(`trajectory_hybrid:237,313`, `yolo_hybrid:407`, `gemini`); **DEFAULT stays `output/`**; golden
+ stub-E2E byte-identical | M1 | ? default-OFF + Launch Report on flip | ? | ? | ? |
117     | M4 | `DeviceContext` + WASB **CPU-fallback**; unified healthcheck/device wiring | M1 |
safe (cuda default unchanged) | ? | ? | ? |
118     | M5 | **truly-local profile** end-to-end (preset + startup checks); first committed
**runlog** (truly-local sample run) | M2,M3,M4 | owner real-GPU | ? | ? | ? |
119     | M6 | Cloud Run GPU **dispatch shim** (control plane ? job via storage ? re-claim) +
**containerized worker image** (ffmpeg+CUDA+WASB+fonts); mocked-dispatch tests | M5 | ? owner
(GCP spend) | ? | ? | ? |
120     | M7 | **cloud-serving profile** e2e on L4: upload`<2GB` + gdrive/bucket ? detect+stitch
on Cloud Run ? reel via signed URL; **DEPLOY_REPORT** (posterity, sample runs + contact sheet) |
M6 | ? owner (spend) | ? | ? | ? |
121     | M8 | Per-job **cost ledger** (v_ migration: `owner_id, gpu_active_seconds,
wall_seconds, bytes_out, cost_usd, cost_breakdown_json`) + **pricebook** + aggregation +
`/api/.../cost` | M7 | ? owner (billing) | ? | ? | ? |
122     | M9 | **Edge auth** (Cf-Access extraction) + per-tenant scoping on `user_id`;
DPA/consent gate wiring | M7 | ? owner (privacy/counsel) | ? | ? | ? |
123     | M10 | `byo_worker` / zero-infra-BYO ? **design-only, DEFERRED** | M8,M9 | ? explicit
owner approval | ? | ? | ? |
124     | **M11** | **Data-flywheel harness (D12):** on a consented adult job, **capture** the
upload + Gemini reel/rally output ? **reliably store** under the per-video workspace ? **shadow-
eval** owned-vs-Gemini (dual-eval-set / `experiment.compare`) ? **promote** good ones to the
golden TRAINING set (human-reviewed labels via the annotation flow). Closes the loop so the
owned model climbs to the D11 floor. Reuses `data_pipeline/` + `backend/eval/` + the consent
gate (M9). | M7,M9 | ? owner (consent/privacy) | ? | ? | ? |
125     | **M12** | **gdrive-api` (OAuth) StorageBackend (D14 Northstar):** cloud reads the
source from + writes the stitched reel **back to the user's Google Drive next to the source**
(per-user OAuth consent) ? the mobile/cloud-native Drive round-trip. Distinct from the local
mount backend; slots into the `StorageBackend` ABC. | M7,M9 | ? owner (OAuth/consent) | ? | ? | ? |
126
127     **Testing strategy is a first-class deliverable** ?
[`TESTING_STRATEGY.md`](TESTING_STRATEGY.md). **Run logs / sample runs for posterity** ?
[`runlogs/`](runlogs/).
128
129     ## 8. Open questions ? owner *(blocking-gates vs nice-to-knows)*
130     **? The 5 gating questions are LOCKED (owner, 2026-06-21) ? see S2 D7?D12:**
131     1. ? **Cold-start ? D7** (scale-to-zero, accept cold-start).
132     3. ? **Pricebook ? D8** (versioned DB table; USD internal / INR display).
133     4. ? **Edge auth ? D9** (Cloudflare Access; verify Cf-Access JWT, lock ingress to

```

```

proxy).
134     5. ? **DPA/consent ? D10** (free-tier data-for-reels trade; adults-only; derived-data;
counsel-gated to M9).
135     6. ? **Quality floor ? D11** + the **data-flywheel harness D12 / M11** (Gemini-ON until
the owned model clears the floor; capture?store?eval?goldenize meanwhile).
136
137     **Still open (non-gating, smaller):**
138     2. **Stitch placement:** always co-located in the detect job (simplest, one metered
unit) vs split to a cheap CPU Cloud Run service when bursty? *Default co-located unless metrics
say otherwise.*
139     7. ? **RESOLVED ? D14:** cloud needs a **`gdrive-api` (OAuth) StorageBackend** (read the
source from + write the reel back to the user's Drive) ? the Northstar's Drive round-trip ?
distinct from the local **mount** backend (which stays for truly-local). New milestone **M12**.
140     8. **NVDEC/NVENC** in the Cloud Run image (cut stitch wall-time/cost) vs libx264
(determinism/parity)?
141     9. **Input cap:** is `<2GB` a hard client reject, or accept to the 4GB upload cap and
route larger files elsewhere?
142     10. **Directory name:** `COMPUTE_DECOUPLED_SERVING/` (chosen) vs keep
`CLOUD_RUN_GPU_SERVING/` ? confirm.
143
144     ## 9. Definition of done
145     - [ ] M1?M4 land (config + input + output-base + device layer), each green on golden
cross-OS + experiment no-regression.
146     - [ ] **Profile-parity test green** ? truly-local vs cloud-mock produce byte-identical
windows + segment ranges.
147     - [ ] **truly-local** (M5) and **cloud-serving** (M7) each have a committed **runlog /
DEPLOY_REPORT** with a real sample run.
148     - [ ] Cost ledger (M8) makes GPU + **stitch** + egress cost visible per job; edge auth
(M9) gates multi-tenant.
149     - [ ] **Data-flywheel harness (M11)** captures?stores?shadow-evals?goldenizes consented
adult jobs; the owned model is on a measured path to the D11 floor.
150     - [ ] ? flip Status DONE, archive to `docs/archives/past_projects/`.
151
152     ## 10. References
153     **ADR lineage:** ADR-005/008/009/010/011/012/013 ? **ADR-014**
(`docs/archives/decisions/DECISIONS.md`). **Strategy:** `PLATFORM_ARCHITECTURE.md`,
`HOSTING_PLAN.md`, `VIDEO_LOCALITY_MODEL.md`, `DESKTOP_APP_PLAN.md`,
`LOCAL_APPLIANCE_ARCHITECTURE.md` (khelsutra). **Archived predecessor:**
`docs/archives/past_projects/DECOUPLED_COMPUTE/` (esp. `05-COST-ATTRIBUTION.md`, `07-COMPUTE-
ABSTRACTION.md`, `06-RISKS-AND-OPEN-QUESTIONS.md`, `DEPLOY_REPORT_GCP_2026-06-20.md`). **Code
anchors `[verified]`:** `backend/pipeline/detectors/base.py:164`, `tracknet_runner.py:281`,
`stitcher/compiler.py:303`, `segmenters/trajectory_hybrid.py:65,237`, `backend/storage/*`,
`backend/worker/`__main__.py:88`, `backend/api/{uploads,job_worker}.py`,
`backend/infrastructure/database.py`. **Design provenance:** workflow `wf_56274ff0` (4 code-
mappers ? synthesis; 58 findings).
154
155     ---
156     ### Changelog
157     - 2026-06-21 ? created from the design workflow output; M0 in progress.
158

```

4. user (2026-06-21T07:24:01.421Z)

```

1     """Deployment-profile presets + environment auto-detection (COMPUTE_DECOUPLED_SERVING
M1).
2
3     A *deployment profile* is a single switch that applies a coherent BUNDLE of existing
4     config knobs, so the truly-local / cloud-serving / cpu-mock modes don't drift apart one
5     knob at a time (the design in ``docs/COMPUTE_DECOUPLED_SERVING/README.md`` §4.3).
6
7     Precedence (lowest ? highest), realised by the loader::
8
9         built-in model defaults < profile preset < config.json < env < worker --set
10
11     This module is **pure data + pure functions** ? no IO, and (critically) **no torch

```

```

import at
12     module load**. The loader applies a preset ONLY when a profile is EXPLICITLY selected
13     (`config.json` ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); with no
14     profile selected the loader is byte-identical to the pre-M1 behaviour (`profile=""` ==
15     today). Auto-detect only ever **suggests** a profile ? it never auto-applies one (D15:
cloud
16     spend is never entered silently).
17     ""
18
19     from __future__ import annotations
20
21     import copy
22     import os
23     from typing import Any, Mapping, Optional
24
25     # Canonical enum values, defined ONCE here. ``models.DeploymentConfig``'s ``Literal``
mirrors
26     # these and ``tests/test_config_deployment.py`` pins parity so the two can never drift.
27     PROFILES: tuple[str, ...] = ("truly-local", "cloud-serving", "cpu-mock")
28     COMPUTE_TARGETS: tuple[str, ...] = ("local_gpu", "cloud_run", "cpu_mock")
29
30     # Each profile maps 1:1 to its canonical compute_target. Used to apply the preset and to
warn
31     # when config.json explicitly overrides compute_target into a state inconsistent with
the profile.
32     COMPUTE_TARGET_FOR_PROFILE: dict[str, str] = {
33         "truly-local": "local_gpu",
34         "cloud-serving": "cloud_run",
35         "cpu-mock": "cpu_mock",
36     }
37
38     # A preset is a partial, nested override dict over the model defaults. It touches ONLY
knobs
39     # that EXIST today ? M1 is pure/additive: the new
``deployment.{profile,compute_target}`` plus
40     # ``indexing.{detector_impl,skip_ai_handoff}`` and ``storage.backend``. Later milestones
extend
41     # these bundles ? ``input_backend`` (M2) and the configurable output/workspace base (M3,
42     # Launch-Report-gated) ? deliberately NOT set here so M1 never pre-empts a gated flip.
43     PROFILE_PRESETS: dict[str, dict[str, Any]] = {
44         "truly-local": {
45             "deployment": {"profile": "truly-local", "compute_target": "local_gpu"},
46             # detector_impl stays "auto" (the runner picks WSL vs native by env on the GPU
box).
47             # Zero-cloud / privacy-first ? run fully local with no Gemini handoff.
48             "indexing": {"skip_ai_handoff": True},
49             "storage": {"backend": "local"},
50         },
51         "cloud-serving": {
52             "deployment": {"profile": "cloud-serving", "compute_target": "cloud_run"},
53             # Linux L4 (cuda); Gemini handoff ON until the owned model clears the #172 floor
(D11).
54             "indexing": {"detector_impl": "native", "skip_ai_handoff": False},
55             "storage": {"backend": "gcs"},
56         },
57         "cpu-mock": {
58             "deployment": {"profile": "cpu-mock", "compute_target": "cpu_mock"},
59             # Deterministic CPU mock ? no GPU/WSL, no provider/API. The CI / regression
profile.
60             "indexing": {"detector_impl": "stub", "skip_ai_handoff": True},
61             "storage": {"backend": "local"},
62         },
63     }
64
65

```



```

66     def is_known_profile(profile: str) -> bool:
67         """True for a non-empty, recognised profile name."""
68         return bool(profile) and profile in PROFILE_PRESETS
69
70
71     def preset_for(profile: str) -> dict[str, Any]:
72         """Return a deep COPY of the preset bundle for ``profile`` (so callers can't mutate
the
73         module-level table), or ``{}`` for ``""`` / an unknown profile."""
74         return copy.deepcopy(PROFILE_PRESETS.get(profile, {}))
75
76
77     def _truthy(val: Optional[str]) -> bool:
78         """Common env-var truthiness (``1/true/yes/on``, case/space-insensitive)."""
79         if val is None:
80             return False
81         return val.strip().lower() in {"1", "true", "yes", "on"}
82
83
84     def _cuda_available() -> bool:
85         """Best-effort CUDA probe. torch is heavy + optional, so it is imported lazily here
and
86         ONLY when a caller opts into ``probe_cuda``; any import/runtime failure means 'no
GPU'."""
87         try:
88             # Intentionally lazy + local: keeps the config-load path torch-free (torch is
heavy
89             # and optional). Only reached when a caller opts into probe_cuda.
90             import torch
91
92             return bool(torch.cuda.is_available())
93         except Exception:
94             return False
95
96
97     def suggest_profile(
98         env: Optional[Mapping[str, str]] = None, *, probe_cuda: bool = False
99     ) -> str:
100         """SUGGEST (never auto-apply ? D15) a deployment profile from the environment.
101
102         Signals, in priority order:
103
104         - ``K_SERVICE`` set (Cloud Run)           ? ``"cloud-serving"``
105         - ``CI`` truthy                           ? ``"cpu-mock"``
106         - ``probe_cuda`` and a CUDA GPU present ? ``"truly-local"``
107         - otherwise                               ? ``""`` (no strong suggestion)
108
109         ``probe_cuda`` defaults to ``False`` so the common config-load path stays cheap and
never
110         imports torch; a CLI ``--detect`` or cloud-serving startup that wants the GPU probe
passes
111         ``probe_cuda=True``. The result is advisory only: the caller decides whether to act
on it.
112
113         """
114         e = os.environ if env is None else env
115         if e.get("K_SERVICE"):
116             return "cloud-serving"
117         if _truthy(e.get("CI")):
118             return "cpu-mock"
119         if probe_cuda and _cuda_available():
120             return "truly-local"
121         return ""

```

```

1      """Configuration loader and saver.
2
3      Provides a single place to load validated AppConfig with built-in defaults
4      merged with optional file overrides (config.json).
5
6      This replaces scattered json.load + .get() patterns throughout the codebase.
7      """
8
9      from __future__ import annotations
10
11     import json
12     import logging
13     import os
14     from pathlib import Path
15     from typing import Any
16
17     from pydantic import BaseModel
18
19     from .models import AppConfig
20     from .presets import (
21         COMPUTE_TARGET_FOR_PROFILE,
22         PROFILE_PRESETS,
23         is_known_profile,
24         preset_for,
25         suggest_profile,
26     )
27
28     logger = logging.getLogger(__name__)
29
30     DEFAULT_CONFIG_PATH = Path("config.json")
31
32
33     def _deep_merge(base: dict[str, Any], override: dict[str, Any]) -> dict[str, Any]:
34         """Recursively overlay ``override`` onto ``base`` (override wins; dicts merge,
35         scalars
36         replace). Used ONLY for the profile-preset precedence path so a config.json sub-key
37         overrides a preset sub-key while sibling preset/default sub-keys survive ? the no-
38         profile
39         path keeps the original shallow ``{**defaults, **file_data}`` merge unchanged."""
40         out = dict(base)
41         for key, val in override.items():
42             if isinstance(val, dict) and isinstance(out.get(key), dict):
43                 out[key] = _deep_merge(out[key], val)
44             else:
45                 out[key] = val
46         return out
47
48     def _resolve_explicit_profile(file_data: dict[str, Any]) -> str:
49         """The EXPLICITLY chosen deployment profile, env winning over config.json. ``""`` if
50         none
51         is set. Auto-detect is deliberately NOT consulted here ? it only suggests (D15), so
52         a bare
53         environment never silently selects a profile (and never silently spends on
54         cloud)."""
55         env_profile = os.environ.get("DEPLOYMENT_PROFILE")
56         if env_profile and env_profile.strip():
57             return env_profile.strip()
58         dep = file_data.get("deployment")
59         if isinstance(dep, dict):
60             prof = dep.get("profile")
61             if isinstance(prof, str) and prof.strip():
62                 return prof.strip()
63         return ""

```



```

119         file_data = {}
120
121     if file_data:
122         unknown = _unknown_key_paths(file_data, AppConfig)
123         if unknown:
124             logger.warning(
125                 "[CONFIG WARNING] %s contains keys the typed config does not "
126                 "recognize (typo'd or removed knobs ? top-level extras are kept "
127                 "but unread; unknown section keys are silently dropped): %s",
128                 cfg_path, ", ".join(sorted(unknown)),
129             )
130
131     # Precedence: built-in defaults < profile preset < config.json (< env/--set, applied
at
132     # consumption sites downstream). A *deployment profile* is opt-in ? it is applied
ONLY when
133     # explicitly selected (DEPLOYMENT_PROFILE env or config.json deployment.profile).
With no
134     # profile the merge is byte-identical to the pre-M1 loader, so the default path is
unchanged.
135     defaults = _get_builtin_defaults()
136     profile = _resolve_explicit_profile(file_data)
137
138     if profile and not is_known_profile(profile):
139         # An unrecognised EXPLICIT profile is loud, not silent (the loader's "dropped
config"
140         # ethos): warn and apply no preset. A bad value coming from config.json
additionally
141         # hits the typed Literal at model_validate below (a hard, clear error); a bad
value
142         # coming only from the env override is ignored here with this warning.
143         logger.warning(
144             "[CONFIG] unrecognized deployment profile %r (valid: %s); applying no
preset.",
145             profile, ", ".join(PROFILE_PRESETS),
146         )
147         profile = ""
148
149     if profile:
150         merged = _deep_merge(_deep_merge(defaults, preset_for(profile)), file_data)
151         # Record the profile actually applied ? the env may have chosen it over
config.json,
152         # so re-stamp it onto the merged deployment section to keep the record honest.
153         merged.setdefault("deployment", {})
154         if isinstance(merged["deployment"], dict):
155             merged["deployment"]["profile"] = profile
156             ct = merged["deployment"].get("compute_target")
157             canonical = COMPUTE_TARGET_FOR_PROFILE.get(profile)
158             if ct and canonical and ct != canonical:
159                 logger.warning(
160                     "[CONFIG] deployment profile %r normally uses compute_target=%r but
"
161                     "config.json overrides it to %r ? honoring the explicit override.",
162                     profile, canonical, ct,
163                 )
164             logger.info("[CONFIG] deployment profile ACTIVE: %s (compute_target=%s)",
165                         profile, merged["deployment"].get("compute_target"))
166     else:
167         # No profile ? the original shallow merge (Pydantic re-defaults nested
sections).
168         merged = {**defaults, **file_data}
169         suggestion = suggest_profile() # cheap, env-only (no torch); SUGGESTS only
(D15)
170         if suggestion:
171             logger.debug(

```

```

172             "[CONFIG] no deployment.profile set (manual knobs). Environment suggests
%r ? "
173             "set deployment.profile or DEPLOYMENT_PROFILE to apply (auto-detect
never "
174             "auto-applies ? D15).", suggestion,
175         )
176
177     return AppConfig.model_validate(merged)
178
179
180 def save_default_config(path: str | Path | None = None) -> Path:
181     """Write a fresh default config.json.
182
183     Used by scripts/doctor.py diagnostics and the --setup CLI path.
184     Overwrites if the file exists.
185     """
186     cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
187     cfg = AppConfig()
188     cfg_path.parent.mkdir(parents=True, exist_ok=True)
189     cfg_path.write_text(cfg.model_dump_json(indent=4), encoding="utf-8")
190     return cfg_path
191
192
193 # Backwards-compatible alias used by some older call sites during transition.
194 def get_default_config() -> dict[str, Any]:
195     """Return defaults as a plain dict (for transitional use only)."""
196     return _get_builtin_defaults()
197

```

6. user (2026-06-21T07:24:02.285Z)

```

1     """Tests for the deployment-profile config (COMPUTE_DECOUPLED_SERVING M1).
2
3     What this pins:
4     - The default profile (") applies NO preset ? the loaded config is byte-identical to
the
5     pre-M1 default, so the core is unaffected (the project's "default-OFF, zero core
change").
6     - Each profile applies its coherent preset bundle (compute_target + the existing
7     detector_impl / skip_ai_handoff / storage.backend knobs).
8     - Precedence: built-in defaults < profile preset < config.json (a file sub-key overrides
a
9     preset sub-key while sibling preset/default sub-keys survive ? i.e. real deep-merge
layering).
10    - D15: auto-detect SUGGESTS but never auto-applies (a bare cloud-like env does NOT
silently
11    select a profile / spend on cloud).
12    - The model's Literal enum stays in parity with backend/config/presets.PROFILE_PRESETS.
13    """
14
15    import json
16    import logging
17    import os
18
19    import pytest
20    from pydantic import ValidationError
21
22    from backend.config import AppConfig, DeploymentConfig, load_config, save_default_config
23    from backend.config.presets import (
24        COMPUTE_TARGET_FOR_PROFILE,
25        COMPUTE_TARGETS,
26        PROFILE_PRESETS,
27        PROFILES,
28        is_known_profile,
29        preset_for,

```

```

30         suggest_profile,
31     )
32
33     NONEXISTENT = "/nonexistent/path/12345/deployment_config.json"
34
35
36     def _write(tmp_path, data) -> str:
37         p = tmp_path / "cfg.json"
38         p.write_text(json.dumps(data))
39         return str(p)
40
41
42     # -----
43     # Defaults / additivity ? the default path must be unchanged.
44     # -----
45     def test_deployment_defaults_are_empty_and_inert():
46         cfg = load_config(NONEXISTENT)
47         assert isinstance(cfg.deployment, DeploymentConfig)
48         assert cfg.deployment.profile == ""
49         assert cfg.deployment.compute_target == ""
50
51
52     def test_default_path_applies_no_preset(monkeypatch):
53         """With no profile selected, the existing knobs keep their built-in defaults (no
54 preset
55 mutates detector_impl / skip_ai_handoff / storage.backend)."""
56         monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
57         cfg = load_config(NONEXISTENT)
58         assert cfg.indexing.detector_impl == "auto"
59         assert cfg.indexing.skip_ai_handoff is False
60         assert cfg.storage.backend == "local"
61
62     def test_default_config_dump_only_adds_deployment_section():
63         """Adding the section must not perturb any other default ? the dump equals a fresh
64 AppConfig()'s dump, and the new section is present with empty values."""
65         dumped = load_config(NONEXISTENT).model_dump(mode="json")
66         assert dumped == AppConfig().model_dump(mode="json")
67         assert dumped["deployment"] == {"profile": "", "compute_target": ""}
68
69
70     # -----
71     # Parity: the model Literal must match the presets table (no drift).
72     # -----
73     def test_profile_literal_matches_presets():
74         import typing
75
76         literal_profiles =
77 set(typing.get_args(DeploymentConfig.model_fields["profile"].annotation))
78         assert literal_profiles == {"", *PROFILES}
79         assert set(PROFILE_PRESETS) == set(PROFILES)
80         assert set(COMPUTE_TARGET_FOR_PROFILE) == set(PROFILES)
81
82     def test_compute_target_literal_matches_presets():
83         import typing
84
85         literal_targets =
86 set(typing.get_args(DeploymentConfig.model_fields["compute_target"].annotation))
87         assert literal_targets == {"", *COMPUTE_TARGETS}
88         assert set(COMPUTE_TARGET_FOR_PROFILE.values()) == set(COMPUTE_TARGETS)
89
90     def test_each_preset_only_touches_existing_knobs():
91         """A preset must only set knobs that exist on the typed models ? otherwise it would

```

```

be a
92     silent no-op. Validates the merged result against AppConfig for every profile."""
93     for profile in PROFILES:
94         # A preset's value set must validate cleanly when merged onto defaults.
95         merged = {**AppConfig().model_dump(mode="json")}
96         for section, kv in preset_for(profile).items():
97             merged.setdefault(section, {})
98             if isinstance(merged[section], dict):
99                 merged[section].update(kv)
100         AppConfig.model_validate(merged) # raises if a preset key/value is invalid
101
102
103     # -----
104     # Per-profile preset application.
105     # -----
106     def test_cpu_mock_profile(tmp_path):
107         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "cpu-mock"}}))
108         assert cfg.deployment.profile == "cpu-mock"
109         assert cfg.deployment.compute_target == "cpu_mock"
110         assert cfg.indexing.detector_impl == "stub"
111         assert cfg.indexing.skip_ai_handoff is True
112         assert cfg.storage.backend == "local"
113
114
115     def test_cloud_serving_profile(tmp_path):
116         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "cloud-serving"}}))
117         assert cfg.deployment.profile == "cloud-serving"
118         assert cfg.deployment.compute_target == "cloud_run"
119         assert cfg.indexing.detector_impl == "native"
120         assert cfg.indexing.skip_ai_handoff is False # Gemini-ON until the owned model
clears the floor
121         assert cfg.storage.backend == "gcs"
122
123
124     def test_truly_local_profile(tmp_path):
125         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "truly-local"}}))
126         assert cfg.deployment.profile == "truly-local"
127         assert cfg.deployment.compute_target == "local_gpu"
128         # detector_impl is intentionally left at the production default (env picks WSL vs
native).
129         assert cfg.indexing.detector_impl == "auto"
130         assert cfg.indexing.skip_ai_handoff is True # zero-cloud
131         assert cfg.storage.backend == "local"
132
133
134     # -----
135     # Precedence: defaults < preset < config.json (deep-merge layering).
136     # -----
137     def test_config_json_overrides_preset_subkey(tmp_path):
138         """A file sub-key beats the preset's sub-key, but a sibling preset sub-key in the
SAME
139         section survives (this is the deep-merge layering ? a shallow merge would wipe
it)."""
140         cfg = load_config(_write(tmp_path, {
141             "deployment": {"profile": "cpu-mock"},
142             "indexing": {"detector_impl": "native"}, # override one preset knob...
143         }))
144         assert cfg.indexing.detector_impl == "native" # file wins
145         assert cfg.indexing.skip_ai_handoff is True # ...preset sibling SURVIVES (deep
merge)
146
147
148     def test_preset_does_not_wipe_unrelated_defaults(tmp_path):
149         """Selecting a profile + setting an unrelated knob keeps the rest of the section at
its

```

```

150     built-in default (the preset and the file both layer onto the full defaults)."""
151     cfg = load_config(_write(tmp_path, {
152         "deployment": {"profile": "cpu-mock"},
153         "indexing": {"ai_max_workers": 2},
154     }))
155     assert cfg.indexing.detector_impl == "stub" # preset
156     assert cfg.indexing.skip_ai_handoff is True # preset
157     assert cfg.indexing.ai_max_workers == 2 # file
158     assert cfg.indexing.default_segmenter == "yolo_hybrid" # untouched default
159     assert cfg.indexing.max_window_duration == 6.0 # untouched default
160
161
162     def test_config_json_can_override_compute_target_with_warning(tmp_path, caplog):
163         cfg_path = _write(tmp_path, {
164             "deployment": {"profile": "cpu-mock", "compute_target": "cloud_run"},
165         })
166         with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
167             cfg = load_config(cfg_path)
168             assert cfg.deployment.profile == "cpu-mock"
169             assert cfg.deployment.compute_target == "cloud_run" # explicit override honored
170         assert any("overrides it to 'cloud_run'" in r.message for r in caplog.records)
171
172
173     # -----
174     # Env selection + precedence over config.json.
175     # -----
176     def test_env_selects_profile(tmp_path, monkeypatch):
177         monkeypatch.setenv("DEPLOYMENT_PROFILE", "cpu-mock")
178         cfg = load_config(_write(tmp_path, {}))
179         assert cfg.deployment.profile == "cpu-mock"
180         assert cfg.indexing.detector_impl == "stub"
181
182
183     def test_env_profile_overrides_config_json_profile(tmp_path, monkeypatch):
184         """env wins over config.json for the profile CHOICE, and the recorded profile
185         matches the
186         preset actually applied (not the file's value)."""
187         monkeypatch.setenv("DEPLOYMENT_PROFILE", "cpu-mock")
188         cfg = load_config(_write(tmp_path, {"deployment": {"profile": "cloud-serving"}}))
189         assert cfg.deployment.profile == "cpu-mock" # env value recorded
190         assert cfg.deployment.compute_target == "cpu_mock" # cpu-mock preset applied
191         assert cfg.indexing.detector_impl == "stub"
192
193     def test_blank_env_profile_is_ignored(tmp_path, monkeypatch):
194         monkeypatch.setenv("DEPLOYMENT_PROFILE", " ")
195         cfg = load_config(_write(tmp_path, {}))
196         assert cfg.deployment.profile == ""
197         assert cfg.indexing.detector_impl == "auto"
198
199
200     def test_unknown_env_profile_warns_and_applies_no_preset(tmp_path, monkeypatch, caplog):
201         monkeypatch.setenv("DEPLOYMENT_PROFILE", "banana")
202         with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
203             cfg = load_config(_write(tmp_path, {}))
204             assert cfg.deployment.profile == "" # not applied
205             assert cfg.indexing.detector_impl == "auto" # no preset
206             assert any("unrecognized deployment profile 'banana'" in r.message for r in
207                 caplog.records)
208
209     def test_unknown_profile_in_config_json_raises(tmp_path, monkeypatch):
210         """A bad VALUE in config.json hits the typed Literal ? a loud, hard error (same
211         class as a

```



```

211         bad `sport`), not a silent drop.""
212     monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
213     with pytest.raises(ValidationError):
214         load_config(_write(tmp_path, {"deployment": {"profile": "banana"}}))
215
216
217     # -----
218     # D15 ? auto-detect SUGGESTS but never AUTO-APPLIES.
219     # -----
220     def test_cloud_env_does_not_auto_apply_profile(tmp_path, monkeypatch):
221         """K_SERVICE (Cloud Run) makes suggest_profile() recommend cloud-serving, but with
no
222         explicit profile the loaded config stays fully default ? NO silent cloud
selection."""
223         monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
224         monkeypatch.setenv("K_SERVICE", "khelesutra-serving")
225         cfg = load_config(_write(tmp_path, {}))
226         assert cfg.deployment.profile == ""          # NOT auto-applied
227         assert cfg.deployment.compute_target == ""
228         assert cfg.indexing.detector_impl == "auto"  # default, not "native"
229         assert cfg.storage.backend == "local"        # default, not "gcs"
230
231
232     def test_active_profile_is_logged(tmp_path, caplog):
233         with caplog.at_level(logging.INFO, logger="backend.config.loader"):
234             load_config(_write(tmp_path, {"deployment": {"profile": "truly-local"}}))
235             assert any("deployment profile ACTIVE: truly-local" in r.message for r in
caplog.records)
236
237
238     def test_no_active_profile_info_when_unset(tmp_path, caplog, monkeypatch):
239         monkeypatch.delenv("DEPLOYMENT_PROFILE", raising=False)
240         with caplog.at_level(logging.INFO, logger="backend.config.loader"):
241             load_config(_write(tmp_path, {}))
242             assert not any("profile ACTIVE" in r.message for r in caplog.records)
243
244
245     # -----
246     # suggest_profile() ? pure env-only logic + the opt-in cuda probe.
247     # -----
248     def test_suggest_profile_cloud_run():
249         assert suggest_profile({"K_SERVICE": "svc"}) == "cloud-serving"
250
251
252     def test_suggest_profile_ci():
253         assert suggest_profile({"CI": "true"}) == "cpu-mock"
254         assert suggest_profile({"CI": "1"}) == "cpu-mock"
255         assert suggest_profile({"CI": "yes"}) == "cpu-mock"
256
257
258     def test_suggest_profile_ci_falsy_or_absent():
259         assert suggest_profile({}) == ""
260         assert suggest_profile({"CI": "false"}) == ""
261         assert suggest_profile({"CI": "0"}) == ""
262         assert suggest_profile({"CI": ""}) == ""
263
264
265     def test_suggest_profile_k_service_beats_ci():
266         assert suggest_profile({"K_SERVICE": "svc", "CI": "true"}) == "cloud-serving"
267
268
269     def test_suggest_profile_no_cuda_probe_by_default(monkeypatch):
270         """probe_cuda defaults False ? no torch import, no truly-local suggestion even with
a GPU."""
271         import backend.config.presets as presets

```

```

272
273 monkeypatch.setattr(presets, "_cuda_available", lambda: True)
274 assert suggest_profile({}) == "" # probe off ? no suggestion
275 assert suggest_profile({}, probe_cuda=True) == "truly-local"
276
277
278 def test_suggest_profile_cuda_probe_false(monkeypatch):
279     import backend.config.presets as presets
280
281     monkeypatch.setattr(presets, "_cuda_available", lambda: False)
282     assert suggest_profile({}, probe_cuda=True) == ""
283
284
285 def test_cuda_available_never_raises(monkeypatch):
286     """Even a broken/missing torch must degrade to False, never propagate."""
287     import builtins
288
289     import backend.config.presets as presets
290
291     real_import = builtins.__import__
292
293     def boom(name, *a, **k):
294         if name == "torch":
295             raise RuntimeError("no torch")
296         return real_import(name, *a, **k)
297
298     monkeypatch.setattr(builtins, "__import__", boom)
299     assert presets._cuda_available() is False
300
301
302 # -----
303 # Helpers + plumbing.
304 # -----
305 def test_is_known_profile():
306     assert is_known_profile("cpu-mock") is True
307     assert is_known_profile("") is False
308     assert is_known_profile("banana") is False
309
310
311 def test_preset_for_returns_copy():
312     a = preset_for("cpu-mock")
313     a["indexing"]["detector_impl"] = "mutated"
314     b = preset_for("cpu-mock")
315     assert b["indexing"]["detector_impl"] == "stub" # module table not mutated
316     assert preset_for("") == {}
317     assert preset_for("banana") == {}
318
319
320 def test_save_default_config_includes_deployment(tmp_path):
321     p = save_default_config(str(tmp_path / "out.json"))
322     saved = json.loads(p.read_text())
323     assert saved["deployment"] == {"profile": "", "compute_target": ""}
324     cfg = load_config(str(p))
325     assert cfg.deployment.profile == ""
326
327
328 def test_repo_config_json_has_no_profile():
329     """The shipped config.json must not silently opt into a profile (default path stays
today's
330     behaviour for the real deployment)."""
331     repo_root = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
332     cfg_path = os.path.join(repo_root, "config.json")
333     if not os.path.exists(cfg_path):
334         pytest.skip("no repo config.json")
335     cfg = load_config(cfg_path)

```

```
336         assert cfg.deployment.profile == ""
337
```

7. assistant (2026-06-21T07:24:05.157Z)

Now let me look at the models.py and __init__.py to understand the full contract, plus the prior test_config.py to understand the existing contract.

8. user (2026-06-21T07:24:06.119Z)

```
1      """Pydantic models for the application configuration.
2
3      This module defines the canonical, validated, typed configuration for the
4      Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6      The models support graceful migration from old config.json files via extra="allow".
7      """
8
9      from __future__ import annotations
10
11     from typing import Any, List, Literal, Optional
12
13     from pydantic import BaseModel, Field, field_validator, model_validator
14
15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0
29         max_scene_cuts_per_minute: float = 0.5
30         relevance: ValidationRelevanceConfig =
31         Field(default_factory=ValidationRelevanceConfig)
32
33     class TrackNetConfig(BaseModel):
34         wsl_distro: str = "Ubuntu"
35         repo_dir: str = "~/models/TrackNetV4"
36         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
37         conda_env: str = "TrackNetV4"
38         weights_path: str = ""
39         queue_length: int = 5
40         stage_in_wsl: bool = True
41         wsl_stage_dir: str = "~/clips"
42         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
43         velocity_threshold: float = 0.02
44         # Cache hygiene: after a SUCCESSFUL run the staged video copy (and, for WASB, the
45         # decoded frame cache) are deleted automatically ? they are pure intermediates,
46         # regenerable from the source, and the dominant disk consumer (~GBs/video). Set
47         # true only for debugging. On FAILURE they are always kept so a re-run can resume.
48         keep_frames: bool = False
49         # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB.
50         # 0 = disabled.
51         min_free_gb: float = 0.0
52
53
54     class WasbIndexConfig(BaseModel):
```

```

55     """Typed config for the live WASB shuttle substrate (indexing.wasb).
56
57     Mirrors backend/pipeline/detectors/wasb_runner.py::WasbConfig.from_indexing_cfg so
58     these knobs actually take effect ? previously this whole block was silently dropped
59     because nested config sections defaulted to extra='ignore'.
60     """
61     wsl_distro: str = "Ubuntu"
62     repo_dir: str = "~/models/WASB-SBDT"
63     conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
64     conda_env: str = "wasb"
65     weights_path: str = "~/models/WASB-
SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
66     sport: str = "badminton"
67     wsl_stage_dir: str = "~/clips"
68     timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
69     velocity_threshold: float = 0.02
70     # --- Linux-native runner only (detector_impl='native'); ignored by the WSL runner.
71     ---
72     # Env vars override these (WASB_DEVICE / WASB_PYTHON) for a config-free GPU box.
73     device: str = "cuda" # torch device for the native runner: cuda | mps | cpu
74     python_bin: str = "python" # the `wasb` conda env python (invoked by path, no
`conda activate`)
75     # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
76     # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates,
regenerable
77     # from the source. Set true only for debugging. On FAILURE they are kept so a re-run
resumes.
78     keep_frames: bool = False
79     # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB. 0 =
disabled.
80     min_free_gb: float = 0.0
81     # FAST PATH ? decode the video on the fly and feed frames straight to the detector,
82     # skipping the slow per-frame PNG extraction that dominates a long clip (~46k PNGs
for a
83     # 12-min 1080p60 video, which overran the 30-min timeout). Output is bit-identical
to the
84     # PNG path (PNG is lossless) ? VERIFIED on a real GPU 2026-06-20 (native stream ==
WSL disk,
85     # 1194/1195 frames byte-identical; native run-to-run byte-identical /
deterministic).
86     # Tri-state:
87     # None = per-runner default ? the WSL runner keeps DISK extraction
(unchanged Windows
88     # behaviour); the Linux-native runner (GPU box) defaults to STREAMING
(the perf win).
89     # True/False = force on/off for BOTH runners (e.g. set False for a container cv2
can't seek).
90     stream_video: Optional[bool] = None
91
92     class FusionConfig(BaseModel):
93         """Config for the multi-signal `fusion_hybrid` segmenter (MULTI_SIGNAL_FUSION_PLAN
P2).
94
95         Every default reproduces control behaviour: the fusion segmenter persists the
96         shuttle net-crossing count (Gate A signal) and ? only when ``cue_flags['person']``
is
97         true ? the person-cue features, but it does NOT gate the served handoff unless a
98         threshold is raised (``min_crossings``/``min_players`` default 0 = OFF). The court
mask
99         is optional: ``court_polygon=None`` ? Gate B counts every detection (player-count-
only);
100         supplied ? off-court persons (spectators / adjacent courts) are excluded.
101         """
102         # Which trajectory substrate seeds the windows (shuttle stays primary).

```

```

103     substrate: Literal["wasb", "tracknet"] = "wasb"
104     # Windowing velocity threshold for the fusion family (the engine reads
105     # indexing.<family>.velocity_threshold). Default matches the substrates' 0.02; set
106     # equal to indexing.wasb.velocity_threshold to A/B fusion against a tuned wasb run.
107     velocity_threshold: float = 0.02
108     # Per-cue enable flags, e.g. {"person": true}. Person cue is OFF unless explicitly
109     # enabled ? so the default path never imports torch / touches the GPU.
110     cue_flags: dict[str, bool] = Field(default_factory=dict)
111     # Served Gate A: drop windows whose shuttle net-crossings < this from the AI
handoff.
112     # 0 = OFF (no gating; persisted candidates are always the full superset for offline
re-scoring).
113     min_crossings: int = Field(default=0, ge=0)
114     # Served Gate B: when the person cue is on, drop windows with median in-court
players
115     # < this. 0 = OFF.
116     min_players: int = Field(default=0, ge=0)
117     # Optional court mask as normalised 0-1 [[x, y], ...] polygon. None = no mask.
118     court_polygon: Optional[list[list[float]]] = None
119     # Net line for the person two-sided-motion split (person features only ? the shuttle
120     # net-crossing gate keeps rally_gate's camera-agnostic auto-estimate).
121     net_axis: Literal["x", "y"] = "y"
122     net_position: float = Field(default=0.5, ge=0.0, le=1.0)
123
124
125     class IndexingConfig(BaseModel):
126         default_segmenter: str = "yolo_hybrid"
127         use_gpu: bool = True
128         # Detector backend for the trajectory hybrids (wasb/tracknet). "auto" = the real
129         # WSL runner (default, production); "stub" = a deterministic CPU mock (no GPU/WSL)
130         # so the whole pipeline is regression-testable on a CPU box / CI; "native" = the
131         # Linux-native WASB runner (no WSL/conda-activate; a GPU box ? M3.5). The
132         # RALLY_DETECTOR_IMPL env var overrides this. See
pipeline/detectors/{stub,native_wasb}_runner.py.
133         detector_impl: str = "auto"
134         motion_threshold: float = 0.08
135         min_segment_duration: float = 3.0
136         dead_time_merge_gap: float = 2.0
137         chunk_duration: float = 90.0
138         chunk_overlap: float = 10.0
139         detector_model: str = "fasterrcnn_resnet50_fpn"
140         detector_score_threshold: float = 0.5
141         yolo_velocity_threshold: float = 0.15
142         yolo_frame_skip: int = 3
143         yolo_tracker: str = "bytetrack.yaml"
144         yolo_prefetch_workers: int = 2
145         min_action_window_duration: float = 2.0
146         # BOUNDED WINDOWING (over-merge fix W1, RALLY_QUALITY_RESEARCH.md §6). 0 = OFF; > 0
= a window
147         # longer than this many seconds is recursively split at its largest internal
inactivity gap (?
148         # ``window_min_split_gap`` s ? the most likely rally boundary), so one window can't
swallow
149         # several rallies + the dead time between them. Never merges, only cuts (recall
can't drop).
150         # *** R1 MILESTONE DEFAULT-ON (cap=6): the smallest safe local-windowing flip.
Measured n=13
151         # golden, R2 tolF1: baseline?milestone (cap=6 + R4 below) ?+0.222, sign 12/0,
p=0.0005; cap=6
152         # beats cap=12 ?+0.110, sign 12/0; net over-seg guard PASS (mean merge_rate
0.651?0.161). ***
153         max_window_duration: float = 6.0
154         window_min_split_gap: float = 0.5 # min internal gap (s) that qualifies as a split
point
155         # HARD-CAP fallback for W1: when True, an over-long window with NO internal

```

```

inactivity gap
156      # (a continuously-active trajectory ? e.g. the real-footage mahadevpura-1 174s blob)
is
157      # chopped at its midpoint until ? max_window_duration, making the cap a true upper
bound.
158      # Only cuts (recall can't drop). OFF by default until measured/promoted.
159      window_hard_cap: bool = False
160      # R4 rally-END inactivity trim (s, 0 = OFF). After a rally the shuttle keeps moving
slowly
161      # (picked up / walked) above velocity_thresh, so the window over-extends its END
(the measured
162      # |?end| gap vs golden). Trim the post-rally tail back to the last frame whose
trailing
163      # window stays dense with FAST motion (velocity > inactivity_rally_thresh). Only
cuts.
164      # *** R1 MILESTONE DEFAULT-ON (1.5/0.20/0.30): on top of cap=6, R4 adds a small but
significant
165      # end-precision gain ? ?tolF1 +0.040, sign 9/1/3, p=0.022, |?end| 4.96?3.31s (n=13).
The W1 cap
166      # is the dominant lever; R4 is the marginal increment. See QUALITY_ITERATIONS.md
"R1". ***
167      window_inactivity_end: float = 1.5
168      inactivity_rally_thresh: float = 0.20
169      inactivity_min_density: float = 0.30
170      # R5 blob-fallback (default OFF). LAST-RESORT splitter for the continuously-active
blob: after
171      # the W1 gap-split AND the R4 inactivity trim have each had their chance, a group
STILL longer
172      # than window_blob_factor × max_window_duration (no internal gap, no rally-end
signal ? e.g.
173      # mahadevpura-1's ~65s residual ? 11× a 6s cap) is midpoint-chopped to ? the cap as
a FORCED cut
174      # (logged + flagged on the window) so the degenerate single-window outcome is
avoided and the
175      # clip is surfaced as needing rally-STATE cues, not a bigger cap. Differs from
window_hard_cap
176      # (which chops BEFORE R4 and over-segments normal clips): blob-fallback runs AFTER
R4 and only
177      # force-cuts the genuine blob (the factor gate). Only cuts (recall can't drop).
178      window_blob_fallback: bool = False
179      # Magnitude gate that keeps R5 surgical: only a group longer than this ×
max_window_duration is a
180      # blob (a normal long rally ~1-2× the cap is left alone; the mahadevpura-1 residual
is ~11×).
181      # Default 4.0 is do-no-harm on the golden corpus (rescues only the continuously-
active
182      # clips, every other clip a no-op within noise) ? see docs/QUALITY_ITERATIONS.md
"R5".
183      window_blob_factor: float = 4.0
184      log_yolo_candidates: bool = True
185      store_yolo_candidates: bool = True
186      ai_handoff_padding: float = 2.0
187      ai_max_workers: int = 5
188      # Width (px) the gemini segmenter re-encodes chunks to before upload. Stream-copied
189      # chunks of a high-bitrate (4K) source blow past the provider upload cap
190      # (backend/providers/gemini.py::MAX_UPLOAD_MB) and are refused ? observed live as
191      # "0 segments / success". Re-encoding also makes chunk boundaries frame-accurate
192      # (stream copy cuts at the nearest keyframe). 0 = legacy stream-copy at source res.
193      # Matches gemini_refine's --clip-scale-width default.
194      ai_chunk_scale_width: int = 1280
195      # FULLY-LOCAL mode (default OFF): when true, the trajectory-hybrid segmenters
196      # (wasb_hybrid / tracknet_hybrid / fusion_hybrid) skip the Phase-3 AI handoff
entirely and
197      # emit their local candidate windows AS the rallies ? no provider, no API key,
nothing

```

```

198         # uploaded. Lets the local detector run standalone (e.g. to benchmark it against the
Gemini
199         # path, or to operate with zero cloud dependency). The pure `gemini` segmenter
ignores this.
200         skip_ai_handoff: bool = False
201         # Optional debug knob: trim every video to the first N seconds before processing.
202         debug_duration: Optional[float] = None
203
204         tracknet: TrackNetConfig = Field(default_factory=TrackNetConfig)
205         wasb: WasbIndexConfig = Field(default_factory=WasbIndexConfig)
206         fusion: FusionConfig = Field(default_factory=FusionConfig)
207
208         @field_validator("default_segmenter")
209         @classmethod
210         def segmenter_must_be_registered(cls, v: str) -> str:
211             # Registry check happens at runtime in main/segmenter selection.
212             # Here we just allow any string for forward compatibility.
213             return v
214
215         @model_validator(mode="after")
216         def _chunk_step_must_be_positive(self) -> "IndexingConfig":
217             # The chunked segmenters advance by (chunk_duration - chunk_overlap); a non-
positive
218             # step makes `while t < duration: t += step` loop forever, growing memory before
any
219             # GPU work. Reject the bad config at load time instead.
220             if self.chunk_overlap >= self.chunk_duration:
221                 raise ValueError(
222                     f"indexing.chunk_overlap ({self.chunk_overlap}) must be < chunk_duration
"
223                     f"({self.chunk_duration}); otherwise chunking never advances."
224                 )
225             return self
226
227
228     class OutputConfig(BaseModel):
229         default_highlights_name: str = "highlights.mp4"
230         db_name: str = "sports_indexer.db"
231         # Directory where the DB, highlight reels, and processed clips are written.
232         # The CLI's --output-dir overrides this at startup (see backend/main.py::main).
233         output_dir: str = "output"
234
235
236     class WatermarkConfig(BaseModel):
237         """Branding overlay burned into each highlight clip during the per-clip re-encode
238         (backend/pipeline/stitcher/compiler.py). **On by default** ? set `enabled: false`
239         (under `stitching.watermark` in config.json) to turn it off. The text is fully
240         configurable. The concat stage stays stream-copy (-c copy)."""
241         enabled: bool = True
242         text: str = "Generated by Khelsutra.guru"
243         logo_path: str = "" # optional transparent PNG overlay; "" = no logo
244         font_path: str = "" # optional .ttf; "" = use the fontconfig `font` family
below
245         font: str = "Sans" # fontconfig family used when font_path is empty
246         font_size_ratio: float = Field(default=0.035, gt=0.0, le=0.5) # text height as a
fraction of frame height
247         opacity: float = Field(default=0.85, ge=0.0, le=1.0)
248         box: bool = True # faint backing box behind the text for legibility
249         margin: int = Field(default=24, ge=0) # px inset from the chosen corner
250         position: Literal["bottom-right", "bottom-left", "top-right", "top-left"] = "bottom-
right"
251
252
253     class StitchingConfig(BaseModel):
254         begin_padding: float = 0.5

```

```

255     end_padding: float = 0.5
256     use_cache: bool = False
257     min_shot_count: int = 1
258     # Long-rally reel filter (seconds; 0 = OFF). Drop detected rally pairs whose
duration
259     # (end ? start) is below this from the highlight reel, BEFORE compiling. A cleaner
260     # reel-selection knob than `min_shot_count`: shot_count is unlabeled / "Unknown" on
the
261     # fully-local no-AI path, whereas duration is derived from the rally's own
start/end. The
262     # local detector over-fires and its false positives are mostly SHORT (motion blips,
263     # background-court fragments), so this keeps the reel to the eye-catching long
rallies.
264     # OFF by default ? the detector output is unchanged, only which rallies reach the
reel.
265     min_rally_duration: float = Field(default=0.0, ge=0.0)
266     watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
267
268
269     class LoggingConfig(BaseModel):
270         """Logging knobs for the run entrypoints (see backend/utils/logging_setup.py)."""
271         # Third-party HTTP/RPC client loggers (httpx, httpcore, urllib3, google-genai,
272         # google.auth, grpc) emit one INFO line per request ? dozens per Gemini run, which
273         # buries our own [rally]/[compile] lines in run.log during manual QA. Default raises
274         # them to WARNING; set true to restore full chatter for request-flow debugging.
275         verbose_http: bool = False
276
277
278     class SecurityConfig(BaseModel):
279         # When set, /api/process_video_path must resolve under this directory (hosted/tenant
280         # mode). Default None preserves the single-user local 'any local file' workflow.
281         allowed_video_root: Optional[str] = None
282
283         # --- Chunked upload (B2, PR-2) ---
284         # Landing zone for browser uploads (assembled from parts). Per-user subdirectories
285         # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
286         # contain upload_dir or /api/process will reject the uploaded paths.
287         upload_dir: str = "output/uploads"
288         # Whole-file cap (default 4 GB) and per-part cap. Parts stay under ~100 MB because
289         # free-tier edge proxies cap request bodies there (HOSTING_PLAN §2).
290         max_upload_mb: int = 4096
291         max_part_mb: int = 95
292         allowed_upload_extensions: List[str] = ["mp4", "mov"]
293
294
295     class StorageConfig(BaseModel):
296         """Pluggable storage backend (docs/DECOUPLED_COMPUTE/). Default ``local`` = today's
297         behaviour (filesystem paths, byte-for-byte). ``s3`` covers AWS + Cloudflare R2 + D0
Spaces
298         + MinIO (via ``endpoint_url``); ``gcs`` = Google Cloud Storage; ``gdrive`` = a
locally MOUNTED
299         Google Drive (``gdrive.mount_root``). Per-backend settings are loose dicts passed to
the constructor ?
300         **secrets are never stored here**, only endpoints / regions / file paths / source-
selectors
301         (boto3/google auth chains supply credentials). See ``backend/storage/``. """
302         backend: Literal["local", "s3", "gcs", "gdrive"] = "local"
303         # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
304         output_root: str = ""
305         # --- Per-video workspace (docs/PER_VIDEO_WORKSPACE.md) ---
306         # One folder per video keyed by the MD5 video_id:
<root>/[workspace_prefix]/<video_id>/...
307         # so a local deploy and a cloud (bucket) deploy share ONE relative layout (only the
root

```



```

308         # differs). DEFAULT-OFF: convergence is phased + flipped via a Launch Report (D5).
When OFF
309         # the legacy flat layout is used unchanged.
310         single_folder_per_video: bool = False
311         # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.
312         workspace_root: str = ""
313         # Cloud namespace prefix under the root so per-video folders sit tidily beside
314         # videos/ labels/ weights/ (e.g. gs://bucket/workspaces/<video_id>/). "" = root
directly.
315         workspace_prefix: str = "workspaces"
316         local: dict[str, Any] = Field(default_factory=dict)
317         s3: dict[str, Any] = Field(default_factory=dict)
318         gcs: dict[str, Any] = Field(default_factory=dict)
319         gdrive: dict[str, Any] = Field(default_factory=dict)
320
321
322     class DeploymentConfig(BaseModel):
323         """Deployment profile ? one switch that selects a coherent bundle of compute +
storage
324         knobs (COMPUTE_DECOUPLED_SERVING M1, ``docs/COMPUTE_DECOUPLED_SERVING/README.md``
§4.3).
325
326         **Pure / additive:** the empty default profile (``""``) applies NO preset, so the
pipeline
327         behaves exactly as before this section existed. A profile is opt-in (``config.json``
328         ``deployment.profile`` or the ``DEPLOYMENT_PROFILE`` env var); auto-detect only ever
329         *suggests* one (D15 ? cloud spend is never entered silently). Nothing in the core
reads
330         these fields yet (M1): selecting a profile works purely by pre-setting the EXISTING
knobs
331         (``indexing.detector_impl`` / ``skip_ai_handoff``, ``storage.backend``) in the
loader.
332
333         Valid values mirror ``backend/config/presets.py`` (parity-pinned by a unit test)."""
334
335         # "" = no profile (today's manual knobs). Others apply presets.PROFILE_PRESETS.
336         profile: Literal["", "truly-local", "cloud-serving", "cpu-mock"] = ""
337         # Where the core runs. Set by the profile preset; overridable in config.json. Inert
in M1
338         # (recorded for observability; consumed by the ComputeBackend seam from M2+).
339         compute_target: Literal["", "local_gpu", "cloud_run", "cpu_mock"] = ""
340
341
342     class AppConfig(BaseModel):
343         """Root configuration object.
344
345         All runtime configuration should be accessed via an instance of this class
346         (or the loader) rather than raw dict lookups.
347         """
348
349         sport: Sport = "badminton"
350         ai_provider: str = "gemini"
351         # Auto-tracking alias: 2.5-flash is legacy now that 3.5-flash is the stable
flagship;
352         # the alias never deprecates underneath us. Eval reports record the run date
alongside.
353         ai_model: str = "gemini-flash-latest"
354
355         validation: ValidationConfig = Field(default_factory=ValidationConfig)
356         indexing: IndexingConfig = Field(default_factory=IndexingConfig)
357         output: OutputConfig = Field(default_factory=OutputConfig)
358         stitching: StitchingConfig = Field(default_factory=StitchingConfig)
359         security: SecurityConfig = Field(default_factory=SecurityConfig)
360         logging: LoggingConfig = Field(default_factory=LoggingConfig)
361         storage: StorageConfig = Field(default_factory=StorageConfig)

```

```

362     deployment: DeploymentConfig = Field(default_factory=DeploymentConfig)
363
364     # Allow extra keys during transition so existing config.json files do not break.
365     model_config = {
366         "extra": "allow",
367         "validate_assignment": True,
368     }
369
370     # Temporary migration shim for legacy dict-style access (config.get("section",
371     {}).get(...)).
372     # Will be removed once all call sites (validators, segmenters, etc.) move to typed
373     # attribute access / the typed config. A nested *section* is returned as a plain
374     dict
375     # via model_dump, so legacy chained access keeps working ? returning the Pydantic
376     # section model instead breaks it (a model has no .get), which bit every validator
377     on
378     # real runs. Scalar leaves are returned as-is.
379     def get(self, key: str, default: Any = None) -> Any:
380         _MISSING = object()
381         val: Any = _MISSING
382         if key in type(self).model_fields or hasattr(self, key):
383             val = getattr(self, key, _MISSING)
384             if val is _MISSING:
385                 # Legacy convenience: find a leaf key inside any nested section. Sections
386                 # are
387                 # discovered dynamically from the model fields, so a future top-level
388                 section
389                 # (e.g. `audio`, `jobs`) is covered automatically without editing a
390                 hardcoded list.
391                 for fname in type(self).model_fields:
392                     section = getattr(self, fname, None)
393                     if isinstance(section, BaseModel) and hasattr(section, key):
394                         val = getattr(section, key)
395                         break
396             if val is _MISSING:
397                 return default
398             if isinstance(val, BaseModel):
399                 return val.model_dump(mode="json")
400             return val
401
402     def __getitem__(self, key: str) -> Any:
403         val = self.get(key)
404         if val is None and key not in type(self).model_fields:
405             # Re-raise like a real dict for unknown top-level
406             raise KeyError(key)
407         return val
408
409     # Convenience type alias for callers that want to annotate
410     Config = AppConfig
411

```

9. user (2026-06-21T07:24:06.576Z)

```

1     """Configuration package.
2
3     Public API:
4     from backend.config import load_config, AppConfig, Config
5
6     This package is the canonical place for all configuration access.
7     """
8
9     from .loader import load_config, save_default_config
10    from .models import ( # re-export common ones
11        AppConfig,

```

```

12         Config,
13         DeploymentConfig,
14         IndexingConfig,
15         StitchingConfig,
16         StorageConfig,
17         ValidationConfig,
18     )
19 from .presets import PROFILE_PRESETS, suggest_profile
20
21 __all__ = [
22     "load_config",
23     "save_default_config",
24     "AppConfig",
25     "Config",
26     "DeploymentConfig",
27     "IndexingConfig",
28     "StitchingConfig",
29     "StorageConfig",
30     "ValidationConfig",
31     "PROFILE_PRESETS",
32     "suggest_profile",
33 ]
34

```

10. user (2026-06-21T07:24:06.819Z)

```

1     """Basic tests for the typed configuration package.
2
3     These ensure:
4     - Defaults are sensible and come from the models.
5     - File overrides work.
6     - Old config files with extra keys are tolerated.
7     - save_default_config produces loadable output.
8     """
9     import json
10    import tempfile
11    import os
12
13
14    from backend.config import load_config, save_default_config, AppConfig
15
16
17    def test_load_default_config():
18        # Force load without the project's config.json so we test the model built-in
19 defaults
20        cfg = load_config("/nonexistent/path/that/does/not/exist/12345/config.json")
21        assert isinstance(cfg, AppConfig)
22        assert cfg.sport == "badminton"
23        assert cfg.indexing.default_segmenter == "yolo_hybrid"
24        assert cfg.validation.min_blur_threshold == 20.0 # updated default from review
25        assert cfg.indexing.ai_max_workers == 5
26
27    def test_logging_config_default_quiets_http():
28        """verbose_http defaults to False so run.log is quiet unless explicitly turned
29 up."""
30        cfg = load_config("/nonexistent/path/12345/config.json")
31        assert cfg.logging.verbose_http is False
32
33    def test_skip_ai_handoff_default_off():
34        """Fully-local mode is opt-in: skip_ai_handoff defaults False so the AI handoff
35 runs."""
36        cfg = load_config("/nonexistent/path/12345/config.json")
37        assert cfg.indexing.skip_ai_handoff is False

```

```

37
38
39 def test_skip_ai_handoff_file_override():
40     with tempfile.TemporaryDirectory() as td:
41         p = os.path.join(td, "skip.json")
42         with open(p, "w") as fh:
43             json.dump({"indexing": {"skip_ai_handoff": True}}, fh)
44         cfg = load_config(p)
45         assert cfg.indexing.skip_ai_handoff is True
46
47
48 def test_logging_config_file_override():
49     with tempfile.TemporaryDirectory() as td:
50         p = os.path.join(td, "logging.json")
51         with open(p, "w") as fh:
52             json.dump({"logging": {"verbose_http": True}}, fh)
53         cfg = load_config(p)
54         assert cfg.logging.verbose_http is True
55
56
57 def test_fusion_config_defaults_reproduce_control():
58     """The fusion block defaults must be no-ops: wasb substrate, no cues, no gating, no
mask.
59
60     Gate A (min_crossings) defaults to 0 = OFF: the #146 served-default flip was
REVERTED
61     (eval?serve ? the +0.074 was on the harness's recall-oriented PROPOSAL windowing; on
the
62     served COARSE windowing Gate A is neutral/negative, ?0.008, dropping real rallies).
See
63     docs/QUALITY_ITERATIONS.md and the served litmus in
tests/test_served_gate_a_regression.py.
64     """
65     cfg = load_config("/nonexistent/path/12345/config.json")
66     f = cfg.indexing.fusion
67     assert f.substrate == "wasb"
68     assert f.cue_flags == {} # person cue OFF
69     assert f.min_crossings == 0 # Gate A OFF (default) ? #146 served flip
reverted
70     assert f.min_players == 0 # Gate B OFF
71     assert f.court_polygon is None # no mask
72     assert f.velocity_threshold == 0.02
73
74
75 def test_fusion_config_gate_a_opt_in_via_min_crossings():
76     """Gate A is opt-in: setting min_crossings explicitly enables it (the eval-only
operating
77     point), while the served default stays OFF (#146 flip reverted)."""
78     with tempfile.TemporaryDirectory() as td:
79         p = os.path.join(td, "optin.json")
80         with open(p, "w") as fh:
81             json.dump({"indexing": {"fusion": {"min_crossings": 2}}}, fh)
82         cfg = load_config(p)
83         assert cfg.indexing.fusion.min_crossings == 2 # Gate A ON only when explicitly
set
84
85
86 def test_served_default_segmenter_is_gemini():
87     """The shipped config.json served default is the prior detector (gemini) ? the #146
flip
88     to fusion_hybrid was REVERTED as a confirmed served regression (eval?serve)."""
89     repo_root = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
90     cfg = load_config(os.path.join(repo_root, "config.json"))
91     assert cfg.indexing.default_segmenter == "gemini"
92     # ...and the served fusion default keeps Gate A OFF (no served gating).

```

```

93         assert cfg.indexing.fusion.min_crossings == 0
94
95
96     def test_fusion_config_file_override():
97         with tempfile.TemporaryDirectory() as td:
98             p = os.path.join(td, "fusion.json")
99             with open(p, "w") as fh:
100                 json.dump({"indexing": {"fusion": {
101                     "substrate": "tracknet", "min_crossings": 2,
102                     "cue_flags": {"person": True}, "court_polygon": [[0.1, 0.1], [0.9, 0.1],
103                     [0.5, 0.9]],
104                     }}}} , fh)
105             cfg = load_config(p)
106             f = cfg.indexing.fusion
107             assert f.substrate == "tracknet"
108             assert f.min_crossings == 2
109             assert f.cue_flags == {"person": True}
110             assert len(f.court_polygon) == 3
111             # An old config without a fusion block still loads (defaults applied).
112             assert cfg.indexing.wasb.velocity_threshold == 0.02
113
114     def test_file_override():
115         with tempfile.TemporaryDirectory() as td:
116             p = os.path.join(td, "override.json")
117             data = {
118                 "indexing": {
119                     "default_segmenter": "gemini",
120                     "ai_max_workers": 3,
121                 },
122                 "validation": {
123                     "min_blur_threshold": 42.0,
124                 },
125             }
126             with open(p, "w") as f:
127                 json.dump(data, f)
128
129             cfg = load_config(p)
130             assert cfg.indexing.default_segmenter == "gemini"
131             assert cfg.indexing.ai_max_workers == 3
132             assert cfg.validation.min_blur_threshold == 42.0
133             # Unspecified keys fall back to defaults
134             assert cfg.sport == "badminton"
135
136
137     def test_extra_keys_tolerated():
138         """Old or future keys in config.json must not break loading."""
139         with tempfile.TemporaryDirectory() as td:
140             p = os.path.join(td, "with_extra.json")
141             data = {
142                 "sport": "badminton",
143                 "indexing": {"default_segmenter": "yolo_hybrid"},
144                 "some_future_key": "ignored_for_now",
145                 "validation": {"relevance": {"court_pixels_min_ratio": 0.1}},
146             }
147             with open(p, "w") as f:
148                 json.dump(data, f)
149
150             cfg = load_config(p)
151             assert cfg.indexing.default_segmenter == "yolo_hybrid"
152             assert cfg.validation.relevance.court_pixels_min_ratio == 0.1
153
154
155     def test_unknown_keys_warn_with_dotted_paths(tmp_path, caplog):
156         """Typo'd keys must be LOUD: top-level extras are kept-but-unread and unknown

```

```

157     section keys are silently dropped, so without this warning a typo just makes
158     the knob a no-op (the audit's 'silently dropped config' class)."""
159     p = tmp_path / "typo.json"
160     p.write_text(json.dumps({
161         "some_future_key": 1,                # top-level extra
162         "indexing": {"velocity_threshold": 0.5}, # typo'd section key
163         "validation": {"relevance": {"unknown_gate": True}}, # doubly nested
164     })))
165     import logging
166     with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
167         cfg = load_config(str(p))
168         assert isinstance(cfg, AppConfig) # still non-fatal
169         joined = " ".join(r.message for r in caplog.records)
170         assert "some_future_key" in joined
171         assert "indexing.velocity_threshold" in joined
172         assert "validation.relevance.unknown_gate" in joined
173
174
175     def test_recognized_config_does_not_warn(tmp_path, caplog):
176         p = tmp_path / "clean.json"
177         p.write_text(json.dumps({
178             "sport": "badminton",
179             "indexing": {"default_segmenter": "gemini", "wasb": {"velocity_threshold":
0.02}},
180         })))
181         import logging
182         with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
183             load_config(str(p))
184         assert not [r for r in caplog.records if "CONFIG WARNING" in r.message]
185
186
187     def test_non_object_config_degrades_to_defaults(tmp_path, caplog):
188         """A config.json that is valid JSON but NOT an object ([], `[1]`, `x`, `42`,
189         `true`,
190         `null`) must degrade to a fully-defaulted AppConfig + a warning ? never crash
191         startup.
192         Regression: previously a truthy non-mapping hit `.items()` (AttributeError) and a
193         falsy one
194         broke the `{**defaults, **file_data}` unpack (TypeError), turning an operator typo
195         into a hard
196         server/CLI startup crash."""
197         import logging
198         default_dump = AppConfig().model_dump(mode="json")
199         for raw in ("[]", "[1, 2]", "x", "42", "true", "null"):
200             p = tmp_path / "bad.json"
201             p.write_text(raw)
202             caplog.clear()
203             with caplog.at_level(logging.WARNING, logger="backend.config.loader"):
204                 cfg = load_config(str(p)) # must not raise
205                 assert isinstance(cfg, AppConfig)
206                 assert cfg.model_dump(mode="json") == default_dump, f"{raw} should yield pure
defaults"
207                 assert any("not a JSON object" in r.message for r in caplog.records), \
208                     f"{raw} should log the non-object warning"
209
210
211     def test_save_default_config_roundtrip():
212         with tempfile.TemporaryDirectory() as td:
213             p = os.path.join(td, "saved.json")
214             saved = save_default_config(p)
215             assert str(saved) == p
216             assert os.path.exists(p)
217
218             cfg = load_config(p)
219             assert cfg.indexing.default_segmenter == "yolo_hybrid"

```

```

216         assert cfg.validation.min_blur_threshold == 20.0
217
218
219 def test_model_is_the_source_of_truth():
220     """Changing a default in the model should be reflected without touching json."""
221     # This is a documentation test more than anything.
222     cfg = AppConfig()
223     assert cfg.indexing.tracknet.velocity_threshold == 0.02
224
225
226 def test_output_dir_field_default_and_override():
227     """output_dir now lives on OutputConfig (was a write-only runtime hack), so
228     /api/compile and the CLI resolve the same value. Default + file override + the
229     runtime mutation that backend/main.py::main performs for --output-dir."""
230     cfg = AppConfig()
231     assert cfg.output.output_dir == "output" # built-in default
232
233     # The CLI override path: main() assigns args.output_dir onto the typed model.
234     cfg.output.output_dir = "/tmp/custom_out"
235     assert cfg.output.output_dir == "/tmp/custom_out"
236
237     # File override is honored on load.
238     with tempfile.TemporaryDirectory() as td:
239         p = os.path.join(td, "out_override.json")
240         with open(p, "w") as f:
241             json.dump({"output": {"output_dir": "reels", "db_name": "x.db"}}, f)
242         loaded = load_config(p)
243         assert loaded.output.output_dir == "reels"
244         assert loaded.output.db_name == "x.db"
245
246 def test_config_shallow_section_replacement():
247     """Positive test for Item 2 (config loader): shallow section replacement.
248
249     File-provided 'indexing' section replaces the defaults 'indexing' section entirely
250     (no deep merge of sub-keys), per loader's {**defaults, **file_data}.
251     """
252     with tempfile.TemporaryDirectory() as td:
253         p = os.path.join(td, "shallow.json")
254         with open(p, "w") as fh:
255             json.dump({"indexing": {"default_segmenter": "motion"}}, fh)
256         cfg = load_config(p)
257         assert cfg.indexing.default_segmenter == "motion"
258         # Pydantic defaults fill the rest of the (replaced) section
259         assert cfg.indexing.ai_max_workers == 5
260
261
262 def test_config_unknown_key_paths_warning_and_tolerate():
263     """Positive test for Item 2: _unknown_key_paths warning for unknown top-level keys,
264     but tolerated (extra='allow'), and unknown section keys dropped silently.
265     """
266     with tempfile.TemporaryDirectory() as td:
267         p = os.path.join(td, "unknown.json")
268         with open(p, "w") as fh:
269             json.dump({
270                 "indexing": {"default_segmenter": "motion"},
271                 "completely_unknown_top": "bar",
272                 "validation": {"unknown_in_section": 99} # dropped
273             }, fh)
274         cfg = load_config(p)
275         assert cfg.indexing.default_segmenter == "motion"
276         # tolerated, no crash
277
278
279 def test_get_returns_dict_for_nested_sections_legacy_compat():
280     """Regression: legacy consumers (validators/segmenters) do

```

```

281     ``config.get("validation", {}).get("min_resolution_width", ...)``. The dict-compat
282     .get must return nested SECTIONS as plain dicts (not Pydantic models), or that
283     chained access raises 'X object has no attribute get' on every real run.
284
285     This is the original strong legacy test (restored per #160 review). It covers the
286     full set of footgun guards: dict sections, nested dicts, scalar leaves with
defaults,
287     missing key fallbacks, and leaf keys discoverable inside sections.
288     """
289     cfg = load_config("/nonexistent/123/config.json")
290     validation = cfg.get("validation", {})
291     assert isinstance(validation, dict)
292     assert validation.get("min_resolution_width") == 1280
293     # nested-of-nested is also dict-accessible
294     assert isinstance(validation.get("relevance"), dict)
295     assert validation["relevance"].get("court_pixels_min_ratio") == 0.05
296
297     # scalar leaves are returned unchanged; a present scalar ignores the default arg;
298     # a missing key falls back to the default
299     assert cfg.get("ai_provider") == "gemini"
300     assert cfg.get("ai_provider", "ignored-default") == "gemini"
301     assert cfg.get("does_not_exist", "fallback") == "fallback"
302
303     # leaf keys inside any section resolve via the dynamic section scan (not a hardcoded
304     # list), so a future top-level section's keys would be discoverable too
305     assert cfg.get("min_resolution_width") == 1280 # lives inside `validation`
306     assert cfg.get("begin_padding") == 0.5         # lives inside `stitching`
307
308     # legacy chained access still works
309     assert cfg.get("indexing", {}).get("default_segmenter") == "yolo_hybrid"
310
311
312     def test_config_roundtrip_and_typed_reaches_pipeline(monkeypatch):
313         """Positive for Item 2: save_default_config + load_config round-trips;
314         typed AppConfig is what reaches pipeline code (e.g. main uses global_config as
AppConfig).
315         """
316         with tempfile.TemporaryDirectory() as td:
317             p = os.path.join(td, "roundtrip.json")
318             save_default_config(p)
319             cfg = load_config(p)
320             assert isinstance(cfg, AppConfig)
321             assert cfg.sport == "badminton"
322
323         # positive that typed is used in pipeline code (global_config is set at startup in
main());
324         # monkeypatch here so the test does not depend on server initialization)
325         import backend.main as main_module
326         from backend.main import get_config
327         monkeypatch.setattr(main_module, "global_config", cfg)
328         assert isinstance(get_config(), AppConfig)
329

```

11. assistant (2026-06-21T07:24:11.032Z)

Now I have a thorough understanding of the contract. Let me run the test suite to confirm the tests pass and check the torch-free guarantee claims more precisely. Let me look closely at the key behaviors.

Let me verify a critical claim: the loader's `suggest\_profile()` call at line 169 uses `os.environ` directly. And I want to check whether the torch-free guarantee is actually tested (the spec specifically calls this out). Let me run the tests first.

12. user (2026-06-21T07:24:14.065Z)

[100%]

.....
31 passed in 0.37s

13. assistant (2026-06-21T07:24:22.066Z)

All 31 tests pass. Now let me analyze the test quality against each loader behavior, per the lens. Let me examine the torch-free guarantee specifically since the spec calls it out as a critical guarantee.

The presets module docstring (line 14-16) and `\_cuda\_available` claim "keeps the config-load path torch-free." Let me check whether any test actually verifies that loading config does NOT import torch.

14. user (2026-06-21T07:24:29.687Z)

```
270:     """probe_cuda defaults False ? no torch import, no truly-local suggestion even with a
GPU."""
286:     """Even a broken/missing torch must degrade to False, never propagate."""
294:         if name == "torch":
295:             raise RuntimeError("no torch")
```

15. user (2026-06-21T07:24:34.869Z)

```
False
---check if torch installed---
torch installed
```

16. assistant (2026-06-21T07:24:40.121Z)

Important finding: torch IS installed in this environment. Now let me analyze the torch-free test claims carefully.

The test `test\_suggest\_profile\_no\_cuda\_probe\_by\_default` (line 269-275) has a docstring claiming "probe_cuda defaults False ? no torch import" but the test only asserts `suggest\_profile({}) == ""`. It does NOT assert that torch was not imported. Since torch is already installed and could be imported elsewhere, this test would pass even if `suggest\_profile` imported torch. Let me verify this is a real gap by checking what would actually catch a torch import in the config-load path.

Let me also examine the env+file conflict and the suggest_profile-uses-real-os.environ behavior, and the compute_target override warning's negative case.

17. user (2026-06-21T07:24:42.837Z)

```
torch in sys.modules after load_config: False
```

18. assistant (2026-06-21T07:24:48.512Z)

Confirmed: `load\_config` does not import torch (because `suggest\_profile()` is called without `probe\_cuda=True` at loader.py:169). But this guarantee is NOT tested ? there is no test asserting torch stays out of `sys.modules` after `load\_config()`. The existing tests would all pass even if `load\_config` imported torch.

Now let me examine the `suggest\_profile()` real-`os.environ` path used by the loader, and the unknown-profile split. Let me check the loader's else-branch suggestion path (line 169) ? it calls `suggest\_profile()` with no env arg, reading the real `os.environ`. Is there a test covering the suggestion log-line in the no-profile path? Let me also scrutinize the deep-merge sibling-survives test more carefully.

19. user (2026-06-21T07:24:54.382Z)

```
default skip_ai_handoff: False
```

20. assistant (2026-06-21T07:25:03.135Z)

Good. The deep-merge-sibling-survives case IS properly tested:

``test_config_json_overrides_preset_subkey`` (line 137-145) asserts ``skip_ai_handoff`` is `True`` while the default is `False``, so a shallow merge would wipe it to `False`` and fail the test. That test passes for the RIGHT reason.

Now let me scrutinize the `compute_target` override warning test for a "passes for wrong reason" issue, and check the ``test_truly_local_profile`` `detector_impl` assertion.

``test_truly_local_profile`` (line 124-131) asserts ``detector_impl == "auto"``. But the truly-local preset does NOT set ``detector_impl`` (presets.py line 47-48, only comment). And the default is `"auto"`. So this assertion ``detector_impl == "auto"`` passes because it's the default anyway ? it would NOT catch a missing preset for `detector_impl`. But that's actually correct here because the preset intentionally doesn't set it. However, it's worth noting this assertion can't distinguish "preset applied" from "preset not applied" for that knob. The ``skip_ai_handoff`` is `True`` and ``compute_target == "local_gpu"`` assertions DO catch preset application, so the test as a whole isn't passing for the wrong reason.

Let me now examine the critical D15 "no-surprise" coverage and whether the `unknown-profile-in-config-json-raises` test is robust, plus check ``test_cloud_env_does_not_auto_apply_profile`` for env leakage.

21. user (2026-06-21T07:25:06.541Z)

CI in env: None

22. assistant (2026-06-21T07:25:14.117Z)

Now let me check one subtle gap: the loader's ``else`` branch (no-profile) at lines 168-175 calls ``suggest_profile()`` and logs at ``logger.debug``. Is there a test that the suggestion is logged in the no-profile path (the D15 "suggest not apply" surfacing)? The test ``test_cloud_env_does_not_auto_apply_profile`` only checks the config values, not that a debug suggestion was logged. That's a minor gap but the D15 surface (logging the suggestion) is part of the "never a surprise ? surface it" contract.

Let me also verify the env+file conflict for `compute_target`. The negative case: when `config.json` does NOT override `compute_target` inconsistently, is the warning NOT emitted? ``test_config_json_can_override_compute_target_with_warning`` checks the warning IS emitted on conflict, but there's no test that the warning is NOT emitted when `compute_target` matches the profile's canonical (a "passes for wrong reason" guard for that warning).

Let me now consolidate and verify the most important finding ? the torch-free guarantee gap ? by confirming the existing tests can't catch a regression.

23. user (2026-06-21T07:25:17.643Z)

Gap confirmed: no test asserts the config-load path is torch-free

24. assistant (2026-06-21T07:25:23.640Z)

I have now thoroughly mapped the test coverage against each loader behavior. Let me do a final verification on two more behaviors the lens asks about: the `unknown-profile file=raise` vs `env=warn` split, and confirm the env+file conflict (env wins) is tested.

The split is well covered:

- ``test_unknown_env_profile_warns_and_applies_no_preset`` (line 200) ? `env=banana` ? warns, no preset, `profile=""`
- ``test_unknown_profile_in_config_json_raises`` (line 209) ? `file=banana` ? `ValidationError`

The env-wins-over-file profile choice is covered by ``test_env_profile_overrides_config_json_profile`` (line 183).

Let me check one last subtle thing ? whether the `unknown-env-profile + valid config.json`

deployment.profile interaction is tested (env=banana but file has profile=cpu-mock: does the env's bad value suppress the file's good one?).

25. user (2026-06-21T07:25:26.335Z)

```
[CONFIG] unrecognized deployment profile 'banana' (valid: truly-local, cloud-serving, cpu-mock);
applying no preset.
profile: 'cpu-mock'
detector_impl: auto
compute_target: ''
```

26. assistant (2026-06-21T07:25:33.246Z)

This reveals a real behavior worth noting: when env=`banana` (invalid) AND config.json has `deployment.profile=cpu-mock` (valid), `\_resolve\_explicit\_profile` returns `banana` (env wins, loader.py:51-53), the loader warns and sets `profile=""` (line 147), so NO preset is applied ? yet the final `cfg.deployment.profile` is `cpu-mock` because the file's raw value survives in `file\_data` and `model\_validate` accepts it. The result is an inconsistent state: `profile="cpu-mock"` but `compute\_target=""` and `detector\_impl="auto"` (preset NOT applied). This is an untested edge interaction. It's arguably a latent bug, but it's an unusual conflict (a bad env var shadowing a good file profile). The lens is test-quality focused, so I'll note this as a missing test case rather than asserting it's a definite bug in M1 scope.

I now have enough to deliver findings. Let me compile the structured output focused on test quality and coverage gaps.

27. user (2026-06-21T07:26:09.789Z)

Structured output provided successfully

28. assistant (2026-06-21T07:26:22.592Z)

I've completed the test-quality review of M1's `tests/test\_config\_deployment.py`. The implementation itself is sound and all 31 tests pass; the findings are about coverage gaps where a regression in loader behavior would NOT be caught.

Summary of findings (5, none blocking):

1. ****Torch-free guarantee untested (major, high)**** ? The most important gap. The spec flags torch-freeness as critical, the code is correct (verified: `load\_config` leaves `torch` out of `sys.modules`), but no test pins it. Since torch is installed here, a regression to eager import would pass all tests. Suggested the exact `sys.modules` assertion.
2. ****No-profile suggestion log-line untested (minor, high)**** ? The D15 "suggest, never apply" surfacing (the DEBUG log telling the operator the env wanted cloud) at loader.py:169-175 has no caplog assertion; deleting it would be silent.
3. ****compute_target warning has no negative case (nit, high)**** ? Only the conflict?warn case is tested; nothing asserts a matching `compute\_target` stays silent, so an unconditional-warning regression passes.
4. ****Bad-env-shadows-good-file conflict untested (minor, high)**** ? Verified live: env=`banana` + file profile=`cpu-mock` yields an inconsistent loaded config (`profile="cpu-mock"` recorded but preset NOT applied: `compute\_target=""`, `detector\_impl="auto"`). This contradicts the "recorded profile == preset applied" invariant the all-valid test pins. No test covers this combination.
5. ****env-vs-file precedence for non-profile knobs absent (nit, medium)**** ? Only the profile *choice* precedence is tested; whether `RALLY\_DETECTOR\_IMPL`/`WASB\_\*` env overrides sit at the top of the chain in M1 is unpinned (may be deferred downstream ? hence medium).

The well-covered behaviors I verified pass for the *right* reason: the deep-merge-sibling-survives case (`test\_config\_json\_overrides\_preset\_subkey` asserts `skip\_ai\_handoff` is True` where the default is `False`, so a shallow merge would fail it), the unknown-profile file=raise vs env=warn split, and per-profile preset application.

Relevant files: `C:\Users\avidu\Projects\badminton-highlight-indexer\tests\test\_config\_deployment.py`, `C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\loader.py`, `C:\Users\avidu\Projects\badminton-highlight-indexer\backend\config\presets.py`.