

claude-windows / dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3\subagents\agent-af0ea96840b5dd1f6.jsonl`
- SHA-256: `e950262fd760450681855663dfb0cfebc196d4eda18447e4a135382469a5fa1c`
- Source modified: `2026-06-21T09:07:53+00:00`
- Imported at: `2026-07-05T16:48:22+00:00`
- project: `subagents`
- session_id: `dc4e53de-0f32-4ae9-8e0e-d6c7d74779f3`
```

Transcript

1. user (2026-06-21T09:04:40.502Z)

Review milestone M4 of COMPUTE_DECOUPLED_SERVING in the repo at C:/Users/avidu/Projects/badminton-highlight-indexer. Run `git diff` (working tree vs HEAD) to see the change.

M4 = "DeviceContext + WASB CPU-fallback; unified healthcheck/device wiring; SAFE (cuda default unchanged)". It adds `backend/pipeline/detectors/device.py::DeviceContext` and wires it into `backend/pipeline/detectors/native\_wasb\_runner.py` so the runner supports a new `device="auto"` value (cuda-if-available-else-cpu) ? the CPU fallback for a no-GPU box ? while keeping explicit `device="cuda"` strict. `indexing.wasb.device` doc updated. New tests: `tests/test\_device\_context.py` + additions to `tests/test\_native\_wasb\_runner.py`.

Verify ONLY these, grounded in code with file:line:

1. SAFETY / "cuda default unchanged": Does an EXPLICIT `device="cuda"` on a GPU-less box still HARD-FAIL the healthcheck (not silently downgrade to cpu)? Is the default still "cuda"? Confirm no behaviour change for any pre-M4 config (cuda/cpu/mps). Is there ANY path where an explicit cuda silently runs on cpu?
2. "auto" NEVER reaches wasb_infer.py: `wasb\_infer.py` would mis-handle a literal "auto" device. Confirm run_predict always resolves auto?cuda/cpu BEFORE building the `--device` argv, on every path (healthcheck-first AND run_predict-without-healthcheck).
3. PROBE CACHING & CORRECTNESS: healthcheck probes cuda and caches it (`self.\_cuda\_available`); run_predict reuses the cache (no 2nd subprocess) but probes lazily if healthcheck was skipped. Is the cache coherent? Could a stale/None cache cause a wrong device? Does `\_probe\_cuda` correctly treat python-missing / timeout / import-error / rc!=0 as "no cuda" (so auto?cpu, never a crash)?
4. EDGE CASES: device="auto" but the torch import itself fails (rc!=0) in healthcheck ? what happens? device="mps" (passthrough)? An unknown device string (e.g. "gpu")? Does the warning fire only on a real fallback?
5. TEST GAPS / OS: are the safety contract (explicit-cuda-strict), the auto?cpu fallback, the auto?cuda, the no-probe-for-explicit, and the cache-reuse all pinned by a test that would FAIL on regression? Any OS-dependent assertion (this runs cross-OS in CI)? Is DeviceContext's frozen/immutability and each property covered?

Report concrete findings with severity (blocker/major/minor/nit) + a fix. A clean bill is valid. Do NOT edit ? review only.

2. assistant (2026-06-21T09:04:42.982Z)

I'll review M4 of COMPUTE_DECOUPLED_SERVING. Let me start by pulling the latest and examining the diff.

3. user (2026-06-21T09:04:44.779Z)

```
backend/config/models.py | 4 +-
backend/pipeline/detectors/native_wasb_runner.py | 55 ++++++-----
tests/test_native_wasb_runner.py | 71 ++++++-----
```

3 files changed, 122 insertions(+), 8 deletions(-)

4. user (2026-06-21T09:04:45.269Z)

```
feat/serving-m3-output-base
aa987cd docs(serving): mark $7 tracker M3 ? + changelog (PR #282)
41b7094 feat(serving): M3 ? configurable output/scratch base (de-hardcode output/, default-OFF)
41d45fd fix(server): drain async jobs under `python -m backend.main` (double-module hang) (#281)
```

5. assistant (2026-06-21T09:04:47.989Z)

I'm on branch `feat/serving-m3-output-base`, but the diff stat shows uncommitted working-tree changes ? that's M4. Let me read the full diff and the relevant files.

6. user (2026-06-21T09:04:48.675Z)

```
diff --git a/backend/config/models.py b/backend/config/models.py
index a31c9da..92ad880 100644
--- a/backend/config/models.py
+++ b/backend/config/models.py
@@ -69,7 +69,9 @@ class WasbIndexConfig(BaseModel):
     velocity_threshold: float = 0.02
     # --- Linux-native runner only (detector_impl='native'); ignored by the WSL runner. ---
     # Env vars override these (WASB_DEVICE / WASB_PYTHON) for a config-free GPU box.
-    device: str = "cuda"          # torch device for the native runner: cuda | mps | cpu
+    # torch device for the native runner. "auto" = cuda-if-available-else-cpu (the M4 CPU
+    fallback
+    # for a no-GPU box); "cuda" stays strict (hard-fails without a GPU ? no silent slow-CPU
+    downgrade).
+    device: str = "cuda"          # auto | cuda | mps | cpu
     python_bin: str = "python"   # the `wasb` conda env python (invoked by path, no `conda
+    activate`)
     # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
     # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates, regenerable
diff --git a/backend/pipeline/detectors/native_wasb_runner.py
b/backend/pipeline/detectors/native_wasb_runner.py
index f8fc1ba..30494a0 100644
--- a/backend/pipeline/detectors/native_wasb_runner.py
+++ b/backend/pipeline/detectors/native_wasb_runner.py
@@ -10,8 +10,9 @@ see the acceptance gate below) ? but WITHOUT the WSL transport:
 * NO ``WslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
   ``wasb`` env's python is invoked **by path**), NO ``/mnt/c`` path translation
   (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
- * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps`` or ``cpu``; threaded
-   into ``wasb_infer.py --device`` (and so into the Hydra ``runner.device`` override).
+ * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps``, ``cpu``, or ``auto``
+   (M4: cuda-if-available-else-cpu ? the CPU fallback for a no-GPU box; explicit ``cuda``
+   stays
+   strict); threaded into ``wasb_infer.py --device`` (and so the Hydra ``runner.device``
+   override).
 * weights / repo / python / scratch come from the **environment** first
   (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
   ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
@@ -40,6 +41,7 @@ from dataclasses import dataclass
from typing import Any, Dict, List, Optional, Tuple

from backend.pipeline.detectors.base import DetectorRunner
+from backend.pipeline.detectors.device import AUTO, DeviceContext
# Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
from backend.pipeline.detectors.tracknet_runner import TrackNetRunner

@@ -62,7 +64,7 @@ class NativeWasbConfig:
    python_bin: str = "python"          # the `wasb` conda env python (invoked by path, no
```

```

activate)
    weights_path: str = "" # REQUIRED (env WASB_WEIGHTS or
indexing.wasb.weights_path)
    sport: str = "badminton"
-   device: str = "cuda" # cuda | mps | cpu
+   device: str = "cuda" # auto | cuda | mps | cpu (auto = cuda-if-available-
else-cpu)
    scratch_dir: str = "" # frame-extraction scratch (env); "" = next to the
output dir
    timeout_sec: int = 1800 # 30 min; a hung GPU call is killed (0 = no timeout)
    keep_frames: bool = False # retain the decoded frame cache after success
(debug only)
@@ -96,6 +98,9 @@ class NativeWasbRunner(DetectorRunner):

    def __init__(self, cfg: NativeWasbConfig):
        self.cfg = cfg
+        # Cached CUDA-availability probe (used to resolve device="auto"). Set by healthcheck;
+        # lazily probed by run_predict if healthcheck was skipped. None = not yet probed.
+        self._cuda_available: Optional[bool] = None

        # --- low-level native exec (the single place tests patch) -----
        def _run(self, argv: List[str], cwd: Optional[str] = None) -> subprocess.CompletedProcess:
@@ -103,6 +108,30 @@ class NativeWasbRunner(DetectorRunner):
        return subprocess.run(argv, cwd=cwd, capture_output=True, text=True,
                               timeout=(self.cfg.timeout_sec or None))

+    # --- device resolution (M4 DeviceContext: device="auto" ? cuda-if-available-else-cpu)
+    -----
+    def _probe_cuda(self) -> bool:
+        """Probe ``torch.cuda.is_available()`` in the wasb env's python (a subprocess). Any
failure
+        (missing python / timeout / import error) is treated as 'no CUDA'."""
+        code = "import torch; print('cuda', torch.cuda.is_available())"
+        try:
+            res = self._run([self.cfg.python_bin, "-c", code])
+        except (FileNotFoundError, subprocess.TimeoutExpired):
+            return False
+        return res.returncode == 0 and "cuda True" in (res.stdout or "")
+
+    def _cuda_is_available(self) -> bool:
+        """CUDA availability, cached across healthcheck + run_predict so we probe at most
once."""
+        if self._cuda_available is None:
+            self._cuda_available = self._probe_cuda()
+        return self._cuda_available
+
+    def _effective_device(self) -> str:
+        """The torch device to actually pass to ``wasb_infer.py``. Only ``auto`` needs the CUDA
+        probe; an explicit cuda/mps/cpu resolves to itself (so ``device='cuda'`` is
unchanged)."""
+        if self.cfg.device == AUTO:
+            return DeviceContext(AUTO, self._cuda_is_available()).effective
+        return self.cfg.device
+
+    def _repo_src(self) -> str:
+        return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")

@@ -143,8 +172,16 @@ class NativeWasbRunner(DetectorRunner):
    out = (res.stdout or "").strip()
    if res.returncode != 0:
        return False, f"{self.cfg.python_bin}` torch import failed: {(res.stderr or
out).strip()}"
-        if self.cfg.device == "cuda" and "cuda True" not in out:
+        # M4: resolve the device. Cache the probe so run_predict reuses it (no second
subprocess).

```

```

+         self._cuda_available = "cuda True" in out
+         ctx = DeviceContext(self.cfg.device, self._cuda_available)
+         if ctx.cuda_required_but_missing:
+             # Unchanged pre-M4 behaviour for an EXPLICIT cuda request: hard-fail, never a
silent
+             # slow-CPU downgrade. Set indexing.wasb.device=auto for a CPU fallback.
+             return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
+         if ctx.fell_back:
+             logger.warning("device=auto: no CUDA GPU detected (%) ? FALLING BACK TO CPU; WASB
"
+
+             "inference will be slow. Set indexing.wasb.device=cuda to require a
GPU.", out)
+         return True, out

    def run_predict(self, video_win_path: str, output_win_dir: str,
@@ -179,13 +216,17 @@ class NativeWasbRunner(DetectorRunner):
        scratch = self.cfg.scratch_dir or output_dir
        frames_dir = os.path.join(scratch, f"{key}_frames")
        argv += ["--frames_out_dir", frames_dir]
+
+        # M4: resolve device="auto" to the effective device (cuda-if-available-else-cpu); an
+        # explicit cuda/mps/cpu passes through unchanged. Never sends "auto" to wasb_infer.py.
+        device = self._effective_device()
        argv += ["--weights", weights,
                "--sport", self.cfg.sport,
-                "--device", self.cfg.device,
+                "--device", device,
+                "--out", out_csv]

-        logger.info("Running native WASB inference (device=%s) on %s ...",
-                    self.cfg.device, os.path.basename(norm))
+        logger.info("Running native WASB inference (device=%s%s) on %s ...",
+                    device, " [auto?fallback]" if self.cfg.device == AUTO and device != "cuda"
else "",
+                    os.path.basename(norm))
+
+        try:
+            res = self._run(argv, cwd=self._repo_src())
+        except FileNotFoundError:
diff --git a/tests/test_native_wasb_runner.py b/tests/test_native_wasb_runner.py
index 22a3abe..2f71f7a 100644
--- a/tests/test_native_wasb_runner.py
+++ b/tests/test_native_wasb_runner.py
@@ -6,6 +6,7 @@ the GPU-box byte-parity gate then relies on: it IS a DetectorRunner, it is NOT a
activate`, no `/mnt` translation) that runs the SAME wasb_infer.py with the SAME output CSV
name as the WSL runner, threads the device through, and cleans the frame cache on success.
"""
+import logging
+import os
+import types
+from unittest.mock import patch
@@ -134,6 +135,76 @@ def test_healthcheck_python_not_found(tmp_path):
    assert not ok and "python" in msg.lower()

+## --- M4: device="auto" CPU fallback (DeviceContext) -----
+def test_healthcheck_auto_uses_cuda_when_available(tmp_path):
+    r = NativeWasbRunner(_ok_cfg(tmp_path, device="auto"))
+    with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda True")):
+        ok, _msg = r.healthcheck()
+        assert ok
+        assert r._cuda_available is True
+        assert r._effective_device() == "cuda" # auto resolves to the GPU when present
+
+
+def test_healthcheck_auto_falls_back_to_cpu_with_warning(tmp_path, caplog):
+    """device=auto on a GPU-less box does NOT hard-fail (unlike device=cuda) ? it resolves to

```

```

cpu
+ and warns loudly. This is the M4 portability win for a no-GPU Linux / cpu-serving box."""
+ r = NativeWasbRunner(_ok_cfg(tmp_path, device="auto"))
+ with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda False")), \
+     caplog.at_level(logging.WARNING,
logger="backend.pipeline.detectors.native_wasb_runner"):
+     ok, _msg = r.healthcheck()
+     assert ok # auto does not require a GPU
+     assert r._effective_device() == "cpu"
+     assert any("FALLING BACK TO CPU" in rec.message for rec in caplog.records)
+
+
+def test_healthcheck_caches_probe_for_run_predict(tmp_path):
+    """healthcheck's CUDA probe is cached so run_predict reuses it (no second subprocess)."""
+    r = NativeWasbRunner(_ok_cfg(tmp_path, device="auto"))
+    with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda False")) as spy:
+        r.healthcheck()
+        assert spy.call_count == 1
+        assert r._effective_device() == "cpu" # uses the cache, no extra probe
+        assert spy.call_count == 1
+
+
+def test_run_predict_auto_threads_cpu_when_no_gpu(tmp_path):
+    """run_predict with device=auto + no GPU: the lazy probe (mocked empty stdout = no cuda)
+    resolves to cpu, and 'auto' is NEVER sent to wasb_infer.py."""
+    _, _out, _rm, argv, _cwd, _key = _drive_success(
+        NativeWasbConfig(weights_path="/m/w.pth", device="auto"), tmp_path)
+    assert argv[argv.index("--device") + 1] == "cpu"
+    assert "auto" not in argv
+
+
+def test_run_predict_auto_threads_cuda_when_gpu(tmp_path):
+    """run_predict with device=auto + a GPU: the probe reports cuda True ? --device cuda."""
+    r = NativeWasbRunner(NativeWasbConfig(weights_path="/m/w.pth", device="auto"))
+    key = "clip_abc123abc123"
+    (tmp_path / f"{key}_wasb.csv").write_text("Frame,Visibility,X,Y\n")
+
+    def _run_side(argv, cwd=None):
+        if "-c" in argv: # the CUDA probe
+            return _proc(0, stdout="cuda True")
+        return _proc(0) # the inference
+
+    with patch.object(r, "_copy_infer_script", return_value=True), \
+        patch.object(r, "_run", side_effect=_run_side) as spy, \
+        patch.object(r, "_rm_rf"):
+        out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id="abc123abc123")
+        assert out is not None
+        infer_argv = next(c.args[0] for c in spy.call_args_list if "wasb_infer.py" in c.args[0])
+        assert infer_argv[infer_argv.index("--device") + 1] == "cuda"
+
+
+def test_run_predict_explicit_cuda_does_not_probe(tmp_path):
+    """An explicit device=cuda resolves to itself with NO probe subprocess ? unchanged
behaviour
+    (a single _run call: the inference)."""
+    r, _out, _rm, argv, _cwd, _key = _drive_success(
+        NativeWasbConfig(weights_path="/m/w.pth", device="cuda"), tmp_path)
+    assert argv[argv.index("--device") + 1] == "cuda"
+    assert r._cuda_available is None # never probed for an explicit device
+
+
+ # --- run_predict -----
def _drive_success(cfg, tmp_path, vid="abc123abc123", **patches):
    """Drive run_predict to success with the subprocess mocked; returns (runner, out, rm_spy,

```

```
argv, cwd)."""
```

7. user (2026-06-21T09:04:49.044Z)

On branch feat/serving-m3-output-base

Your branch is based on 'origin/feat/serving-m3-output-base', but the upstream is gone.

(use "git branch --unset-upstream" to fixup)

Changes not staged for commit:

(use "git add <file>..." to update what will be committed)

(use "git restore <file>..." to discard changes in working directory)

modified: backend/config/models.py

modified: backend/pipeline/detectors/native_wasb_runner.py

modified: tests/test_native_wasb_runner.py

Untracked files:

(use "git add <file>..." to include in what will be committed)

artifacts/

backend/pipeline/detectors/device.py

docs/KHELSUTRA_PRODUCT_OVERVIEW.md

tests/test_device_context.py

no changes added to commit (use "git add" and/or "git commit -a")

8. assistant (2026-06-21T09:04:51.706Z)

The diff doesn't show the new untracked files (`device.py` and `test\_device\_context.py`). Let me read those plus the full `native\_wasb\_runner.py` to ground my review in actual line numbers.

9. user (2026-06-21T09:04:52.190Z)

```
1      """DeviceContext ? resolve the effective torch device for a detector runner.
2
3      COMPUTE_DECOUPLED_SERVING M4 (realizes the ``DeviceContext`` designed in the archived
4      ``07-COMPUTE-ABSTRACTION``). A runner asks for a device by name; this turns the
5      *requested*
6      device + a probed CUDA availability into the *effective* device to actually run on, and
7      says
8      whether a CPU fallback happened.
9
10     The portability gap it closes: the native WASB runner hardcoded ``device='cuda'`` and
11     its
12     healthcheck hard-failed on a box with no GPU, so a CPU-only Linux box (cpu-serving / a
13     no-GPU truly-local install / CI) could not run at all.
14
15     **Safety contract (so "cuda default unchanged" holds):**
16     * ``"auto"`` ? ``"cuda"`` if a GPU is available, else ``"cpu"`` (the fallback). This is
17     opt-in.
18     * ``"cuda"`` ? honored verbatim; an *explicit* cuda request on a GPU-less box still
19     hard-fails
20     elsewhere (no silent slow-CPU downgrade of a config that asked for the GPU).
21     * ``"cpu"`` / ``"mps"`` ? honored verbatim.
22
23     """
24
25     from __future__ import annotations
26
27     from dataclasses import dataclass
28
29     AUTO = "auto"
30     CUDA = "cuda"
31     CPU = "cpu"
32
33     @dataclass(frozen=True)
```

```

29     class DeviceContext:
30         """The requested device + a probed CUDA availability ? the effective device."""
31
32         requested: str          # what config/env asked for: auto | cuda | mps | cpu
33         cuda_available: bool    # probed in the runner's own torch env (not the main
process)
34
35         @property
36         def effective(self) -> str:
37             """The device the runner should actually use. ``auto`` resolves to cuda-if-
available-else-cpu;
38             any explicit device is returned verbatim."""
39             if self.requested == AUTO:
40                 return CUDA if self.cuda_available else CPU
41             return self.requested
42
43         @property
44         def fell_back(self) -> bool:
45             """True when ``auto`` resolved to CPU because no GPU was found (worth a loud
warning)."""
46             return self.requested == AUTO and not self.cuda_available
47
48         @property
49         def cuda_required_but_missing(self) -> bool:
50             """True when an EXPLICIT ``cuda`` request can't be satisfied ? the caller should
hard-fail
51             rather than silently downgrade (preserving the pre-M4 strict behaviour for
``device='cuda'``)."""
52             return self.requested == CUDA and not self.cuda_available
53

```

10. user (2026-06-21T09:04:52.552Z)

```

1     """DeviceContext ? torch device resolution for detector runners
(COMPUTE_DECOUPLED_SERVING M4).
2
3     Pins the safety contract: ``auto`` falls back cuda?cpu (the new portability win), while
an
4     explicit ``cuda`` is honored verbatim and flagged as required-but-missing (the caller
hard-fails,
5     no silent slow-CPU downgrade ? "cuda default unchanged").
6     """
7     from backend.pipeline.detectors.device import DeviceContext
8
9
10    def test_auto_resolves_to_cuda_when_available():
11        ctx = DeviceContext("auto", cuda_available=True)
12        assert ctx.effective == "cuda"
13        assert ctx.fell_back is False
14        assert ctx.cuda_required_but_missing is False
15
16
17    def test_auto_falls_back_to_cpu_when_no_gpu():
18        ctx = DeviceContext("auto", cuda_available=False)
19        assert ctx.effective == "cpu"
20        assert ctx.fell_back is True
21        assert ctx.cuda_required_but_missing is False
22
23
24    def test_explicit_cuda_is_honored_verbatim():
25        assert DeviceContext("cuda", cuda_available=True).effective == "cuda"
26        miss = DeviceContext("cuda", cuda_available=False)
27        assert miss.effective == "cuda" # honored ? NOT silently downgraded
to cpu
28        assert miss.cuda_required_but_missing is True # ?so the caller hard-fails

```

```

29         assert miss.fell_back is False
30
31
32     def test_explicit_cpu_and_mps_pass_through():
33         assert DeviceContext("cpu", cuda_available=False).effective == "cpu"
34         assert DeviceContext("cpu", cuda_available=True).effective == "cpu"
35         assert DeviceContext("mps", cuda_available=False).effective == "mps"
36         for d in ("cpu", "mps"):
37             ctx = DeviceContext(d, cuda_available=False)
38             assert ctx.fell_back is False
39             assert ctx.cuda_required_but_missing is False
40
41
42     def test_is_frozen():
43         import dataclasses
44         import pytest
45         ctx = DeviceContext("auto", cuda_available=True)
46         with pytest.raises(dataclasses.FrozenInstanceError):
47             ctx.requested = "cpu" # type: ignore[misc]
48

```

11. assistant (2026-06-21T09:04:55.728Z)

Now let me read the full `native\_wasb\_runner.py` to see the healthcheck and run_predict in context, especially the ordering of operations and how `healthcheck` is (or isn't) called before `run\_predict`.

12. user (2026-06-21T09:04:56.724Z)

```

1     """Linux-native WASB shuttle-trajectory runner (no WSL).
2
3     The de-risk goal the ``DetectorRunner`` ABC was built for (see ``base.py`` ?
4     "Optional Linux-native path"; ``docs/DECOUPLED_COMPUTE/04-BUILD-AND-TEST-PLAN.md``
5     M3.5).
6     It is the sibling of :class:`~backend.pipeline.detectors.wasb_runner.WasbRunner` for a
7     Linux GPU box (GCP / a SaaS worker): it runs the **same** verified ``wasb_infer.py``
8     core
9     that the WSL adapter runs, so the trajectory CSV is identical (to a sub-pixel float
10    tolerance;
11    see the acceptance gate below) ? but WITHOUT the WSL transport:
12
13    * NO ``WslCommandMixin`` (no ``wsl -d <distro> bash -c``), NO ``conda activate`` (the
14    ``wasb`` env's python is invoked **by path**), NO ``/mnt/c`` path translation
15    (``to_wsl_mnt_path``). Paths are native; the video is already local on the box.
16    * ``device`` is ``cuda`` (default ? the WSL parity path), ``mps``, ``cpu``, or
17    ``auto``
18    (M4: cuda-if-available-else-cpu ? the CPU fallback for a no-GPU box; explicit
19    ``cuda`` stays
20    strict); threaded into ``wasb_infer.py --device`` (and so the Hydra
21    ``runner.device`` override).
22    * weights / repo / python / scratch come from the **environment** first
23    (``WASB_WEIGHTS`` / ``WASB_REPO_DIR`` / ``WASB_PYTHON`` / ``WASB_SCRATCH``), then
24    ``indexing.wasb.*`` config, then a default ? so a box needs no config edits.
25
26    It emits the identical ``Frame,Visibility,X,Y`` CSV (same ``{stem}__vid12_wasb.csv``
27    name as the WSL runner) and reuses the model-agnostic ``parse_trajectory_csv`` /
28    ``trajectory_to_action_windows`` from ``tracknet_runner``, so the trajectory hybrids and
29    the windowing brain are unchanged. Selected via the existing ``detector_impl`` seam
30    (``trajectory_hybrid._resolve_runner`` ? ``_build_native_runner``);
31    ``detector_impl="native"``.
32
33    **Acceptance gate (parity, NOT strict byte-equality):** the SAME clip through the WSL
34    runner
35    (Windows) and this runner (Linux, ``device=cuda``) yields the same trajectory CSV + the
36    same

```



```

28     candidate windows **within a small float tolerance (~1e-3 px)**. cuDNN is non-
deterministic
29     ACROSS GPUs, so a sub-pixel diff is expected on different hardware ? but it is byte-
exact
30     run-to-run on a FIXED GPU. ? Validated 2026-06-20 on an RTX 2070 (via WSL): native
streaming
31     vs the cached WSL-disk CSV matched **1194/1195** frames (the lone diff a ~5e-5 px edge
32     artifact), and two native runs were byte-identical (deterministic). This is a hardware
check
33     (a real GPU + the ``wasb`` env), so the offline, subprocess-mocked suite asserts the
contract,
34     not the numbers.
35     """
36     import os
37     import shutil
38     import logging
39     import subprocess
40     from dataclasses import dataclass
41     from typing import Any, Dict, List, Optional, Tuple
42
43     from backend.pipeline.detectors.base import DetectorRunner
44     from backend.pipeline.detectors.device import AUTO, DeviceContext
45     # Reuse the model-agnostic helpers ? trajectory shape + windowing are identical to
WSL/TrackNet.
46     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
47
48     logger = logging.getLogger(__name__)
49
50     # Path to the inference core we copy into the WASB repo's src/ (so it can import WASB's
51     # detectors/trackers/dataloaders + find configs/), exactly as the WSL adapter does.
52     _INFER_PY = os.path.join(os.path.dirname(os.path.abspath(__file__)), "wasb_infer.py")
53
54
55     @dataclass
56     class NativeWasbConfig:
57         """Resolved settings for a Linux-native WASB run.
58
59         Built via :meth:`from_indexing_cfg` with **environment overrides** taking precedence
60         over ``indexing.wasb.*`` config (the box sets env; no config edits needed).
61         ``device`` defaults to ``cuda`` (the WSL parity path); ``stream_video`` defaults to True (the
perf win).
62         """
63         repo_dir: str = "~/models/WASB-SBDT" # native clone of nttcom/WASB-SBDT
64         python_bin: str = "python" # the `wasb` conda env python (invoked by
path, no activate)
65         weights_path: str = "" # REQUIRED (env WASB_WEIGHTS or
indexing.wasb.weights_path)
66         sport: str = "badminton"
67         device: str = "cuda" # auto | cuda | mps | cpu (auto = cuda-if-
available-else-cpu)
68         scratch_dir: str = "" # frame-extraction scratch (env); "" = next
to the output dir
69         timeout_sec: int = 1800 # 30 min; a hung GPU call is killed (0 = no
timeout)
70         keep_frames: bool = False # retain the decoded frame cache after
success (debug only)
71         # NATIVE DEFAULT = STREAMING: skip the cv2 PNG extraction that dominates a long
clip's
72         # wall-clock (measured ~50s on a 20s clip) ? bit-identical to disk (verified on a
real GPU
73         # 2026-06-20). The GPU box reads from a local NVMe so streaming is the right default
here.
74         stream_video: bool = True
75

```

```

76     @classmethod
77     def from_indexing_cfg(cls, idx_cfg: Dict[str, Any]) -> "NativeWasbConfig":
78         w = (idx_cfg or {}).get("wasb", {}) or {}
79         env = os.environ.get
80         # stream_video is tri-state in config: None/absent => native default (stream);
81         # explicit True/False still wins (e.g. False for a container cv2 can't seek).
82         sv = w.get("stream_video")
83         return cls(
84             repo_dir=env("WASB_REPO_DIR") or w.get("repo_dir", "~/models/WASB-SBDT"),
85             python_bin=env("WASB_PYTHON") or w.get("python_bin", "python"),
86             weights_path=env("WASB_WEIGHTS") or w.get("weights_path", "") or "",
87             sport=w.get("sport", "badminton"),
88             device=env("WASB_DEVICE") or w.get("device", "cuda"),
89             scratch_dir=env("WASB_SCRATCH") or env("RALLY_SCRATCH") or env("TMPDIR") or
90             "",
91             timeout_sec=int(w.get("timeout_sec", 1800)),
92             keep_frames=bool(w.get("keep_frames", False)),
93             stream_video=(True if sv is None else bool(sv)),
94         )
95
96     class NativeWasbRunner(DetectorRunner):
97         """WASB runner that executes natively on a Linux GPU box ? no WSL. See module
98         docstring."""
99         def __init__(self, cfg: NativeWasbConfig):
100             self.cfg = cfg
101             # Cached CUDA-availability probe (used to resolve device="auto"). Set by
102             # healthcheck;
103             # lazily probed by run_predict if healthcheck was skipped. None = not yet
104             # probed.
105             self._cuda_available: Optional[bool] = None
106             # --- low-level native exec (the single place tests patch)
107             -----
108             def _run(self, argv: List[str], cwd: Optional[str] = None) ->
109             subprocess.CompletedProcess:
110                 logger.debug("native exec: %s (cwd=%s)", " ".join(argv), cwd)
111                 return subprocess.run(argv, cwd=cwd, capture_output=True, text=True,
112                                     timeout=(self.cfg.timeout_sec or None))
113
114             # --- device resolution (M4 DeviceContext: device="auto" ? cuda-if-available-else-
115             # cpu) -----
116             def _probe_cuda(self) -> bool:
117                 """Probe ``torch.cuda.is_available()`` in the wasb env's python (a subprocess).
118                 Any failure
119                 (missing python / timeout / import error) is treated as 'no CUDA'."""
120                 code = "import torch; print('cuda', torch.cuda.is_available())"
121                 try:
122                     res = self._run([self.cfg.python_bin, "-c", code])
123                 except (FileNotFoundError, subprocess.TimeoutExpired):
124                     return False
125                 return res.returncode == 0 and "cuda True" in (res.stdout or "")
126
127             def _cuda_is_available(self) -> bool:
128                 """CUDA availability, cached across healthcheck + run_predict so we probe at
129                 most once."""
130                 if self._cuda_available is None:
131                     self._cuda_available = self._probe_cuda()
132                 return self._cuda_available
133
134             def _effective_device(self) -> str:
135                 """The torch device to actually pass to ``wasb_infer.py``. Only ``auto`` needs
136                 the CUDA

```

```

130         probe; an explicit cuda/mps/cpu resolves to itself (so ``device='cuda'`` is
unchanged)."""
131         if self.cfg.device == AUTO:
132             return DeviceContext(AUTO, self._cuda_is_available()).effective
133         return self.cfg.device
134
135     def _repo_src(self) -> str:
136         return os.path.join(os.path.expanduser(self.cfg.repo_dir), "src")
137
138     def _copy_infer_script(self) -> bool:
139         """Copy wasb_infer.py into the WASB repo src/ so it imports WASB modules + finds
configs/."""
140         repo_src = self._repo_src()
141         try:
142             os.makedirs(repo_src, exist_ok=True)
143             shutil.copy(_INFER_PY, os.path.join(repo_src, "wasb_infer.py"))
144             return True
145         except OSError as e:
146             logger.error("Failed to copy wasb_infer.py into %s: %s", repo_src, e)
147             return False
148
149     def _rm_rf(self, path: Optional[str]) -> None:
150         """Best-effort recursive delete of a native path (post-success cleanup; never
raises)."""
151         if path:
152             shutil.rmtree(path, ignore_errors=True)
153
154     def healthcheck(self) -> Tuple[bool, str]:
155         """Verify weights + repo present and the `wasb` python imports torch (and sees a
GPU on cuda)."""
156         if not self.cfg.weights_path:
157             return False, ("WASB weights path is empty ? set WASB_WEIGHTS (or
indexing.wasb.weights_path).")
158         weights = os.path.expanduser(self.cfg.weights_path)
159         if not os.path.exists(weights):
160             return False, f"WASB weights not found at {weights}"
161         repo_src = self._repo_src()
162         if not os.path.isdir(repo_src):
163             return False, (f"WASB-SBDT repo src/ not found at {repo_src} ? clone
nttcom/WASB-SBDT "
164                             f"(set WASB_REPO_DIR / indexing.wasb.repo_dir).")
165         code = "import torch; print('torch', torch.__version__, 'cuda',
torch.cuda.is_available())"
166         try:
167             res = self._run([self.cfg.python_bin, "-c", code])
168         except FileNotFoundError:
169             return False, f"python binary not found: {self.cfg.python_bin!r} (set
WASB_PYTHON)."
170         except subprocess.TimeoutExpired:
171             return False, "native torch healthcheck timed out."
172         out = (res.stdout or "").strip()
173         if res.returncode != 0:
174             return False, f"`{self.cfg.python_bin}` torch import failed: {(res.stderr or
out).strip()}"
175         # M4: resolve the device. Cache the probe so run_predict reuses it (no second
subprocess).
176         self._cuda_available = "cuda True" in out
177         ctx = DeviceContext(self.cfg.device, self._cuda_available)
178         if ctx.cuda_required_but_missing:
179             # Unchanged pre-M4 behaviour for an EXPLICIT cuda request: hard-fail, never
a silent
180             # slow-CPU downgrade. Set indexing.wasb.device=auto for a CPU fallback.
181             return False, f"device=cuda but torch.cuda.is_available() is False ({out})"
182         if ctx.fell_back:
183             logger.warning("device=auto: no CUDA GPU detected (%s) ? FALLING BACK TO

```

```

CPU; WASB "
184             "inference will be slow. Set indexing.wasb.device=cuda to
require a GPU.", out)
185         return True, out
186
187     def run_predict(self, video_win_path: str, output_win_dir: str,
188                     video_id: Optional[str] = None) -> Optional[str]:
189         """Run WASB natively on a video. Returns the path to the trajectory CSV, or
None.
190
191         Output artifacts are namespaced by ``video_id`` (the file MD5) exactly as the
WSL
192         runner does, and the CSV is named ``{stem}__{vid12}_wasb.csv`` ? byte-for-byte
the
193         same downstream contract, so the parity comparison is apples-to-apples.
194         """
195         output_dir = os.path.abspath(output_win_dir)
196         os.makedirs(output_dir, exist_ok=True)
197         # Normalize for a cross-platform basename (tests pass Windows paths on Linux).
198         norm = str(video_win_path).replace("\\", "/")
199         stem, _ext = os.path.splitext(os.path.basename(norm))
200         # Content-derived key: same filename + different bytes -> different key (no
cache reuse).
201         key = f"{stem}__{video_id[:12]}" if video_id else (stem or "wasb")
202         out_csv = os.path.join(output_dir, f"{key}_wasb.csv")
203
204         if not self._copy_infer_script():
205             logger.error("Could not stage wasb_infer.py into the WASB repo src/.")
206             return None
207
208         weights = os.path.expanduser(self.cfg.weights_path)
209         argv = [self.cfg.python_bin, "wasb_infer.py",
210                "--video", os.path.abspath(str(video_win_path))]
211         frames_dir: Optional[str] = None
212         if self.cfg.stream_video:
213             # Fast path: decode in-process, no PNG extraction. Output is bit-identical
to disk.
214             argv.append("--stream-video")
215         else:
216             scratch = self.cfg.scratch_dir or output_dir
217             frames_dir = os.path.join(scratch, f"{key}_frames")
218             argv += ["--frames_out_dir", frames_dir]
219             # M4: resolve device="auto" to the effective device (cuda-if-available-else-
cpu); an
220             # explicit cuda/mps/cpu passes through unchanged. Never sends "auto" to
wasb_infer.py.
221             device = self._effective_device()
222             argv += ["--weights", weights,
223                    "--sport", self.cfg.sport,
224                    "--device", device,
225                    "--out", out_csv]
226
227             logger.info("Running native WASB inference (device=%s%s) on %s ...",
228                         device, " [auto?fallback]" if self.cfg.device == AUTO and device !=
"cuda" else "",
229                         os.path.basename(norm))
230             try:
231                 res = self._run(argv, cwd=self._repo_src())
232             except FileNotFoundError:
233                 logger.error("python binary not found: %r (set WASB_PYTHON).",
self.cfg.python_bin)
234                 return None
235             except subprocess.TimeoutExpired:
236                 logger.error("WASB native inference timed out after %ss.",
self.cfg.timeout_sec)

```

```

237         return None
238     if res.returncode != 0:
239         logger.error("WASB native inference failed:\n%s", (res.stderr or res.stdout
or "")[-2000:])
240         return None
241     if not os.path.exists(out_csv):
242         logger.error("WASB native finished but expected CSV not found at %s",
out_csv)
243         return None
244     logger.info("WASB trajectory CSV: %s", out_csv)
245
246     # Success ? the CSV is the durable output; the decoded frame cache is a pure
intermediate.
247     # On any earlier failure we return above WITHOUT cleaning, so the cache survives
a resume.
248     # (Streaming never materializes frames, so there is nothing to drop on that
path.)
249     if not self.cfg.keep_frames and frames_dir:
250         logger.info("Cache hygiene: removing frame cache for %s "
251                     "(set indexing.wasb.keep_frames=true to retain).", key)
252         self._rm_rf(frames_dir)
253     return out_csv
254
255     # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
256     parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
257     trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)
258

```

13. assistant (2026-06-21T09:05:00.922Z)

Now I have the full picture. Let me verify a few subtleties: the `\_probe\_cuda` stdout parsing vs healthcheck stdout, the `\_drive\_success` helper to confirm what the test asserts, and how `healthcheck` relates to `run\_predict` in production (whether healthcheck is always called first).

14. user (2026-06-21T09:05:02.344Z)

```

200     """An explicit device=cuda resolves to itself with NO probe subprocess ? unchanged
behaviour
201     (a single _run call: the inference)."""
202     r, _out, _rm, argv, _cwd, _key = _drive_success(
203         NativeWasbConfig(weights_path="/m/w.pth", device="cuda"), tmp_path)
204     assert argv[argv.index("--device") + 1] == "cuda"
205     assert r._cuda_available is None          # never probed for an explicit device
206
207
208     # --- run_predict
-----
209     def _drive_success(cfg, tmp_path, vid="abc123abc123", **patches):
210         """Drive run_predict to success with the subprocess mocked; returns (runner, out,
rm_spy, argv, cwd)."""
211         r = NativeWasbRunner(cfg)
212         key = f"clip__{vid[:12]}"
213         (tmp_path / f"{key}_wasb.csv").write_text("Frame,Visibility,X,Y\n") # success = CSV
exists
214         with patch.object(r, "_copy_infer_script", return_value=True), \
215             patch.object(r, "_run", return_value=_proc(0)) as run_spy, \
216             patch.object(r, "_rm_rf") as rm_spy:
217             out = r.run_predict(r"C:\v\clip.mp4", str(tmp_path), video_id=vid)
218             argv = run_spy.call_args[0][0]
219             cwd = run_spy.call_args.kwargs.get("cwd")
220             return r, out, rm_spy, argv, cwd, key
221

```

```

222
223     def test_run_predict_builds_native_argv_no_wsl(tmp_path):
224         # Pin disk mode here (default is now streaming) so this also covers the
225         --frames_out_dir path.
226         _, out, rm, argv, cwd, key = _drive_success(
227             NativeWasbConfig(weights_path="/m/w.pth", device="cuda", stream_video=False),
228             tmp_path)
229         assert out == os.path.join(os.path.abspath(str(tmp_path)), f"{key}_wasb.csv")
230         # No WSL / conda plumbing anywhere in the command.
231         assert "wsl" not in argv
232         assert not any("conda" in a or "/mnt/" in a for a in argv)
233         # It runs the SAME inference core, device-threaded, writing the SAME CSV name as
234         WSL.
235         assert argv[0] == "python" and argv[1] == "wasb_infer.py"
236         assert "--device" in argv and argv[argv.index("--device") + 1] == "cuda"
237         assert "--weights" in argv and argv[argv.index("--weights") + 1] == "/m/w.pth"
238         assert argv[argv.index("--out") + 1] == out
239         # Disk mode extracts frames + cleans them on success.
240         assert "--frames_out_dir" in argv and "--stream-video" not in argv
241         assert cwd.endswith(os.path.join("src"))
242         rm.assert_called_once()
243
244     def test_run_predict_default_is_stream_mode_no_frames_dir(tmp_path):
245         # No stream_video set -> the NATIVE DEFAULT (streaming): no PNG extraction, nothing
246         to clean.
247         _, out, rm, argv, _cwd, _key =
248         _drive_success(NativeWasbConfig(weights_path="/m/w.pth"), tmp_path)
249         assert out is not None
250         assert "--stream-video" in argv and "--frames_out_dir" not in argv
251         rm.assert_not_called() # streaming never materializes a frame cache
252
253     def test_run_predict_frames_dir_uses_scratch_and_is_cleaned(tmp_path):
254         """Frames extract under scratch_dir (the box's big NVMe) when set, and the SAME dir
255         is the
256         one cleaned on success ? a regression that dropped scratch_dir (frames -> small
257         system disk)
258         or diverged the cleanup target would be a disk-fill leak."""
259         scratch = tmp_path / "scratch"
260         r, out, rm, argv, _cwd, key = _drive_success(
261             NativeWasbConfig(weights_path="/m/w.pth", scratch_dir=str(scratch),
262             stream_video=False), tmp_path)
263         expected_frames = os.path.join(str(scratch), f"{key}_frames")
264         assert argv[argv.index("--frames_out_dir") + 1] == expected_frames
265         rm.assert_called_once_with(expected_frames) # cleanup target == extraction target
266
267     def test_run_predict_frames_dir_defaults_next_to_output(tmp_path):
268         _, out, rm, argv, _cwd, key = _drive_success(
269             NativeWasbConfig(weights_path="/m/w.pth", stream_video=False), tmp_path)
270         expected_frames = os.path.join(os.path.abspath(str(tmp_path)), f"{key}_frames")
271         assert argv[argv.index("--frames_out_dir") + 1] == expected_frames
272         rm.assert_called_once_with(expected_frames)
273
274     def test_copy_infer_script_copies_byte_identical(tmp_path):
275         """Unmocked: the native runner stages the SAME wasb_infer.py the WSL runner runs
276         (the
277         byte-parity precondition) into <repo>/src/, byte-for-byte."""
278         import filecmp
279         from backend.pipeline.detectors import native_wasb_runner as nwr
280         r = NativeWasbRunner(NativeWasbConfig(repo_dir=str(tmp_path / "repo")))
281         assert r._copy_infer_script() is True
282         dest = tmp_path / "repo" / "src" / "wasb_infer.py"

```

```

278         assert dest.exists() and filecmp.cmp(nwr._INFER_PY, str(dest), shallow=False)
279
280
281     def test_copy_infer_script_false_on_oserror(tmp_path, monkeypatch):
282         r = NativeWasbRunner(NativeWasbConfig(repo_dir=str(tmp_path / "repo")))
283         monkeypatch.setattr("backend.pipeline.detectors.native_wasb_runner.shutil.copy",
284                             lambda *a, **k: (_ for _ in []).throw(SError("no space")))
285         assert r._copy_infer_script() is False
286
287
288     def test_run_predict_device_threaded_cpu(tmp_path):
289         _, _out, _rm, argv, _cwd, _key = _drive_success(

```

15. user (2026-06-21T09:05:03.535Z)

```

backend\results.py:6:(e.g. missing API key, WSL healthcheck failure, provider error, etc.).
backend\worker\_main_.py:240:    ok, msg = runner.healthcheck()
backend\reporting\spec.yaml:64: # provider error / quota / WSL healthcheck issue rather than a
truly empty video.
backend\reporting\spec.yaml:74:    failed WSL healthcheck ? these all return [] silently. Re-
run with -u for
backend\pipeline\segmenters\base.py:30:    "no segments found" from "something went wrong"
(missing key, healthcheck
backend\pipeline\detectors\base.py:11:- Provide a single place to evolve the healthcheck / error
contract and thread
backend\pipeline\detectors\base.py:21:- healthcheck() -> tuple[bool, str]
backend\pipeline\detectors\base.py:31:2. Implement run_predict and healthcheck (return (bool,
str) tuple).
backend\pipeline\detectors\base.py:41:6. Add unit tests that exercise the contract (isinstance,
healthcheck shape, run
backend\pipeline\detectors\base.py:51:- healthcheck can just check "weights file exists and
torch importable + CUDA?".
backend\pipeline\detectors\base.py:57:The healthcheck message is already logged by the hybrids
on failure. The reporting
backend\pipeline\detectors\base.py:179:    def healthcheck(self) -> Tuple[bool, str]:
backend\pipeline\detectors\device.py:9:healthcheck hard-failed on a box with no GPU, so a CPU-
only Linux box (cpu-serving / a
backend\pipeline\detectors\native_wasb_runner.py:101:    # Cached CUDA-availability probe
(used to resolve device="auto"). Set by healthcheck;
backend\pipeline\detectors\native_wasb_runner.py:102:    # lazily probed by run_predict if
healthcheck was skipped. None = not yet probed.
backend\pipeline\detectors\native_wasb_runner.py:123:    """CUDA availability, cached across
healthcheck + run_predict so we probe at most once."""
backend\pipeline\detectors\native_wasb_runner.py:154:    def healthcheck(self) -> Tuple[bool,
str]:
backend\pipeline\detectors\native_wasb_runner.py:171:    return False, "native torch
healthcheck timed out."
backend\pipeline\segmenters\tracknet_hybrid.py:9:The pipeline engine (healthcheck -> trajectory
-> windows -> AI handoff -> DB)
backend\pipeline\segmenters\trajectory_hybrid.py:7:    healthcheck gate -> P2 trajectory ->
candidate windows (+ persistence)
backend\pipeline\segmenters\trajectory_hybrid.py:227:    # --- Phase 1: env healthcheck
(fail fast with a clear message) ---
backend\pipeline\segmenters\trajectory_hybrid.py:228:    ok, msg = runner.healthcheck()
backend\pipeline\segmenters\wasb_hybrid.py:11:The pipeline engine (healthcheck -> trajectory ->
windows -> AI handoff -> DB)
backend\pipeline\detectors\tracknet_runner.py:97:    def healthcheck(self) -> Tuple[bool, str]:
backend\pipeline\detectors\tracknet_runner.py:112:    return False, "WSL healthcheck
timed out."
backend\pipeline\detectors\stub_runner.py:70:    def healthcheck(self) -> Tuple[bool, str]:
backend\pipeline\detectors\wasb_runner.py:72:    def healthcheck(self) -> Tuple[bool, str]:
backend\pipeline\detectors\wasb_runner.py:82:    return False, "WSL healthcheck timed
out."

```

16. assistant (2026-06-21T09:05:06.252Z)

Let me verify the production call path ? `trajectory\_hybrid.py` around line 227-228 ? to confirm whether healthcheck is always run before `run\_predict` (relevant to cache coherence) and whether a failed healthcheck stops execution.

17. user (2026-06-21T09:05:07.366Z)

```
220         "detector": self.name,
221         "velocity_thresh": velocity_thresh,
222         "merge_gap": merge_gap,
223         "min_action_window_duration": min_window_duration,
224         **runner_params,
225     }
226
227     # --- Phase 1: env healthcheck (fail fast with a clear message) ---
228     ok, msg = runner.healthcheck()
229     if not ok:
230         logger.error("%s detector environment not ready: %s", label, msg)
231         return []
232     logger.info("%s detector env OK: %s", label, msg)
233
234     # --- Phase 2: trajectory -> candidate rally windows ---
235     # Trajectory detectors process the whole video in one pass (they carry
236     # their own frame buffering), so unlike yolo_hybrid we don't pre-chunk.
237     t_start = time.time()
238     out_dir = os.path.join(output_base(self.config), self.family)
239     csv_path = runner.run_predict(video_path, out_dir, video_id=video_id)
240     if not csv_path:
241         logger.error("%s produced no trajectory; aborting.", label)
242         return []
243
244     points = runner.parse_trajectory_csv(csv_path)
245     logger.info("Parsed %d trajectory points (%d visible).",
246               len(points), sum(1 for p in points if p.visible))
247
248     all_candidate_windows = runner.trajectory_to_action_windows(
249         points, fps=fps, frame_width=frame_width, chunk_start=0.0,
250         velocity_thresh=velocity_thresh, merge_gap=merge_gap,
251         min_window_duration=min_window_duration, chunk_index=0,
252         max_window_duration=max_window_duration, min_split_gap=window_min_split_gap,
253         window_hard_cap=window_hard_cap,
254         window_inactivity_end=window_inactivity_end,
255         inactivity_rally_thresh=inactivity_rally_thresh,
256         inactivity_min_density=inactivity_min_density,
257         window_blob_fallback=window_blob_fallback,
258         window_blob_factor=window_blob_factor,
259     )
260     logger.info("Phase 2 (%s) completed in %.1fs. Candidate windows: %d",
261               label, time.time() - t_start, len(all_candidate_windows))
262
263     if log_candidates:
264         for cw in all_candidate_windows:
265             logger.info(f"[{label}] rally
266             { _fmt_ts(cw['start']) } - { _fmt_ts(cw['end']) } "
267             f"({cw['end'] - cw['start']:.1f}s) |
268             frames={cw['active_frame_count']} "
269             f"peak_v={cw['peak_velocity']:.3f}
270             mean_v={cw['mean_velocity']:.3f}")
271
272     # Per-window stats via the enrichment hook ? base = the 4 motion keys; the
273     fusion
274     # segmenter adds net_crossings + person-cue features. Computed once, reused for
275     both
276     # persistence and the handoff gate below.
277     window_stats = [
278         self._enrich_candidate_stats(cw, points, fps, frame_width, video_path,
```



```
idx_cfg)
274             for cw in all_candidate_windows
```

18. assistant (2026-06-21T09:05:11.899Z)

This confirms the production path: `healthcheck()` is called at line 228 and `run_predict()` at line 239 on the `**same runner instance**` ? so in production, the probe cache set in `healthcheck` is always populated and reused by `run_predict`. The lazy probe in `run_predict` is only for callers (tests, direct API) that skip `healthcheck`.

Now let me run the new tests and the full suite to verify green, and check lint/type cleanliness on the changed files.

19. user (2026-06-21T09:05:21.010Z)

```
..... [100%]
43 passed in 5.15s
```

20. assistant (2026-06-21T09:05:22.708Z)

All 43 tests pass. Let me run the full suite and lint/type checks on the changed files.

21. user (2026-06-21T09:06:46.980Z)

```
=====
Full test suite: C:\Users\avidu\AppData\Local\Programs\Python\Python313\python.exe -m pytest
-q
obsreport (optional reporting dep): available
=====
..... [ 6%]
..... [ 12%]
..... [ 19%]
..... [ 25%]
..... [ 32%]
..... [ 38%]
..... [ 44%]
..... [ 51%]
..... [ 57%]
..... [ 64%]
..... [ 70%]
..... [ 76%]
..... [ 83%]
..... [ 89%]
..... [ 96%]
..... [100%]
1125 passed in 80.69s (0:01:20)
=====
RESULT: PASS - all tests green
=====
```

22. user (2026-06-21T09:06:47.810Z)

```
All checks passed!
---MYPY---
Success: no issues found in 2 source files
```

23. assistant (2026-06-21T09:06:53.496Z)

Full suite green (1125 passed), ruff clean, mypy clean. Now let me carefully verify each of the 5 review points. Let me check one subtle correctness concern: the `_probe_cuda` parses `"cuda True"` from `print('cuda', torch.cuda.is_available())`, while `healthcheck` parses from `print('torch', ..., 'cuda', ...)`. Let me confirm both produce a substring `"cuda True"` correctly, and check the mps-on-auto and unknown-device edge cases by tracing the logic precisely.

Let me also verify there's no `mps`-availability concern and check how `wasb\_infer.py` handles the device argument, to confirm the "auto never reaches it" claim matters.

24. user (2026-06-21T09:06:55.129Z)

```
cuda True
---
torch 1.11.0 cuda False
```

25. user (2026-06-21T09:06:56.450Z)

```
56: def _build_cfg(weights: str, sport: str, device: str = "cuda"):
57:     """Compose the WASB eval config via Hydra, overriding model/weights/device.
59:     ``device`` is ``cuda`` (default ? the historical WSL behaviour, byte-parity preserved),
61:     cpu/mps run with no GPU list so the detector places tensors on ``device``.
64:     gpus = "[0]" if device == "cuda" else "[]" # single GPU on cuda; none for cpu/mps
70:         f"runner.device={device}",
138:         frames_in: int, frames_total: int, device: str = "cuda") -> bool:
141:     ``device`` defaults to ``cuda`` so caches written before device-keying (no ``device``
142:     field) stay compatible with a CUDA run ? but a cpu/mps run will NOT reuse a cuda cache
143:     (detections can differ numerically across devices), and vice-versa.
153:         and manifest.get("device", "cuda") == device
376:         device: str = "cuda") -> str:
407:     cfg = _build_cfg(weights, sport, device)
443:         frames_in=frames_in, frames_total=len(frames), device=device,
461:         "device": device,
648:         device: str = "cuda") -> str:
666:         frame_loader=frame_loader, source_key=source_key, device=device)
682:     ap.add_argument("--device", default="cuda", choices=["cuda", "mps", "cpu"],
683:         help="torch device for the detector (default cuda = the WSL parity
path)")
701:         flush_every=args.flush_every, fresh=args.fresh,
device=args.device)
712:         flush_every=args.flush_every, fresh=args.fresh, device=args.device)
```

26. assistant (2026-06-21T09:07:03.528Z)

This confirms a key point: `wasb\_infer.py` line 682-683 declares `--device` with `choices=["cuda", "mps", "cpu"]` ? argparse would **reject** a literal `auto` with exit code 2. So the claim "auto must never reach wasb_infer.py" is real and load-bearing: if `auto` leaked through, the inference subprocess would hard-fail. The M4 code correctly resolves `auto` ? cuda/cpu before building argv on both paths.

Now let me trace the edge cases precisely to confirm point 4. Let me verify the `\_effective\_device` behavior for `mps` and unknown strings, and the torch-import-failure path in healthcheck.

Let me re-examine the healthcheck ordering for the `auto` + torch-import-failure case. At line 173, if `res.returncode != 0` (torch import failed), healthcheck returns `False` early ? before reaching the device-resolution block at line 176. So `\_cuda\_available` stays `None`. Let me trace what happens next if a caller ignores the healthcheck failure and calls `run\_predict` anyway with `device="auto"`.

27. user (2026-06-21T09:07:04.281Z)

```
100
101
102     def test_healthcheck_repo_missing(tmp_path):
103         w = tmp_path / "w.pth.tar"
104         w.write_text("x")
105         cfg = NativeWasbConfig(weights_path=str(w), repo_dir=str(tmp_path / "absent"))
106         ok, msg = NativeWasbRunner(cfg).healthcheck()
107         assert not ok and "repo" in msg.lower()
108
```

```

109
110     def test_healthcheck_torch_import_fails(tmp_path):
111         r = NativeWasbRunner(_ok_cfg(tmp_path))
112         with patch.object(r, "_run", return_value=_proc(1, stderr="ModuleNotFoundError:
torch")):
113             ok, msg = r.healthcheck()
114             assert not ok and "failed" in msg.lower()
115
116
117     def test_healthcheck_cuda_unavailable_when_device_cuda(tmp_path):
118         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cuda"))
119         with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda
False")):
120             ok, msg = r.healthcheck()
121             assert not ok and "cuda" in msg.lower()
122
123
124     def test_healthcheck_cpu_device_ok_without_cuda(tmp_path):
125         r = NativeWasbRunner(_ok_cfg(tmp_path, device="cpu"))
126         with patch.object(r, "_run", return_value=_proc(0, stdout="torch 1.11.0 cuda
False")):
127             ok, msg = r.healthcheck()
128             assert ok # cpu run does not require a GPU
129
130
131     def test_healthcheck_python_not_found(tmp_path):
132         r = NativeWasbRunner(_ok_cfg(tmp_path, python_bin="/no/such/python"))
133         with patch.object(r, "_run", side_effect=FileNotFoundError()):
134             ok, msg = r.healthcheck()
135             assert not ok and "python" in msg.lower()
136
137

```

28. assistant (2026-06-21T09:07:15.273Z)

I now have everything I need to complete the review. Let me trace the remaining edge cases precisely by reasoning through the code (no need to run more):

****Edge case A ? `auto` + torch import fails (rc!=0) in healthcheck:**** healthcheck returns `False` at line 173-174 (import failed), `\_cuda\_available` stays `None`. Production aborts (trajectory_hybrid line 229-231). If a caller skips healthcheck and calls `run\_predict` with `device="auto":` `\_effective\_device` ? `\_cuda\_is\_available` ? `\_probe\_cuda` runs its own subprocess (`import torch; print(...)`). That subprocess also fails (rc!=0 since torch unimportable) ? `\_probe\_cuda` returns `False` ? device resolves to `cpu`. Then the actual inference subprocess will fail too (torch unimportable), and `run\_predict` returns `None` gracefully. No crash. Correct.

****Edge case B ? `mps`:**** `\_effective\_device` line 131: `self.cfg.device == AUTO` is False ? returns `"mps"` verbatim, no probe. healthcheck: `DeviceContext("mps", ...)` ? `cuda\_required\_but\_missing` False, `fell\_back` False ? passes, no warning. Correct passthrough.

****Edge case C ? unknown device e.g. `gpu`:**** `\_effective\_device` returns `gpu` verbatim (not AUTO). healthcheck: `cuda\_required\_but\_missing` False (not "cuda"), `fell\_back` False ? healthcheck passes. Then `wasb\_infer.py --device gpu` hits argparse `choices=["cuda","mps","cpu"]` ? subprocess exits rc=2 ? `run\_predict` returns `None`. This is a pre-existing behavior (M4 didn't change validation of unknown strings), and it fails loudly at inference, not silently. Acceptable ? noting as a minor.

****Edge case D ? warning fires only on real fallback:**** line 182 `if ctx.fell\_back` ? only true when `requested==AUTO` and not `cuda\_available`. cpu/mps/explicit-cuda never set `fell\_back`. Correct.

Let me confirm there are no other consumers of `.device` that could leak `auto`, and check the worker `\_\_main\_\_.py` path that also calls healthcheck.

29. user (2026-06-21T09:07:17.275Z)

```
backend/pipeline/detectors/native_wasb_runner.py:15:    strict); threaded into ``wasb_infer.py
--device`` (and so the Hydra ``runner.device`` override).
backend/pipeline/detectors/native_wasb_runner.py:44:from backend.pipeline.detectors.device
import AUTO, DeviceContext
backend/pipeline/detectors/native_wasb_runner.py:128:    def _effective_device(self) -> str:
backend/pipeline/detectors/native_wasb_runner.py:131:        if self.cfg.device == AUTO:
backend/pipeline/detectors/native_wasb_runner.py:133:            return self.cfg.device
backend/pipeline/detectors/native_wasb_runner.py:177:        ctx =
DeviceContext(self.cfg.device, self._cuda_available)
backend/pipeline/detectors/native_wasb_runner.py:180:            # slow-CPU downgrade. Set
indexing.wasb.device=auto for a CPU fallback.
backend/pipeline/detectors/native_wasb_runner.py:184:            "inference will
be slow. Set indexing.wasb.device=cuda to require a GPU.", out)
backend/pipeline/detectors/native_wasb_runner.py:221:        device = self._effective_device()
backend/pipeline/detectors/native_wasb_runner.py:228:            device, "
[auto?fallback]" if self.cfg.device == AUTO and device != "cuda" else "",
backend/pipeline/segmenters/fusion_hybrid.py:87:
"fusion_substrate": "wasb", "device": cfg.device}
backend/pipeline/segmenters/wasb_hybrid.py:46:
"detector_impl": "native", "device": cfg.device}
```

30. assistant (2026-06-21T09:07:53.584Z)

Lines 87 (fusion_hybrid) and 46 (wasb_hybrid) put `cfg.device` into `runner\_params` for reporting/metadata only ? they don't pass it to argv. The only argv path is `\_effective\_device()` at line 221. Both `--device` argv construction sites (line 224 via `\_effective\_device`) are covered. The `auto` string is fully contained.

Let me do one final check ? confirm `\_cuda\_available` being set to `False` (not `None`) for explicit non-cuda devices in healthcheck doesn't cause an inconsistency, and verify the cache-coherence concern about a `False` cache being reused correctly when device is `auto`.

In healthcheck line 176: `self.\_cuda\_available = "cuda True" in out` is set for **every** device (cuda/cpu/mps/auto) once torch imports successfully. This is a genuine probed value, so the cache is always coherent with reality. For `auto`, run_predict's `\_effective\_device` ? `\_cuda\_is\_available` sees the cached bool (not None) ? no second probe. Correct and coherent.

One subtle thing to verify: is there any scenario where the cache is set to a **stale** value? The probe reflects the box's GPU at healthcheck time; GPU availability doesn't change within a process lifetime, so no staleness risk. Clean.

The review is complete. All checks pass.

Review: M4 (DeviceContext + WASB CPU-fallback) ? COMPUTE_DECOUPLED_SERVING

Verdict: clean bill. No blockers, no majors. Two nits. The "cuda default unchanged" safety contract holds, `auto` never reaches `wasb\_infer.py`, the probe cache is coherent, and the safety/fallback/cache contracts are all pinned by tests that would fail on regression.

Verified locally: full suite **1125 passed** (`python run\_all\_tests.py`), the M4 tests **43 passed**, `ruff check` clean, `mypy` clean on both changed source files.

1. SAFETY / "cuda default unchanged" ? PASS

- Default is still `cuda`: `backend/config/models.py:72` and `native\_wasb\_runner.py:67` both default `device="cuda"`; `from\_indexing\_cfg` falls back to `"cuda"` (`native\_wasb\_runner.py:88`).
- Explicit `cuda` on a GPU-less box still **hard-fails** the healthcheck: `native\_wasb\_runner.py:177-181` ? `DeviceContext.cuda\_required\_but\_missing` (`device.py:49-52`) is True only for `requested=="cuda" and not cuda\_available`, returning the unchanged `device=cuda` but torch.cuda.is_available() is False` error. Pinned by `test_healthcheck_cuda_unavailable_when_device_cuda` and `test_explicit_cuda_is_honored_verbatim`.

- No silent cuda?cpu path: `\_effective\_device` (`native\_wasb\_runner.py:128-133`) only consults the probe when `device == "auto"`; explicit `cuda` returns `"cuda"` verbatim. `DeviceContext.effective` (`device.py:39-41`) likewise returns explicit devices verbatim. cpu/mps unchanged (`test\_healthcheck\_cpu\_device\_ok\_without\_cuda`, `test\_explicit\_cpu\_and\_mps\_pass\_through`).

2. "auto" NEVER reaches wasb_infer.py ? PASS (and load-bearing)

Confirmed `wasb\_infer.py:682` declares `--device ... choices=["cuda","mps","cpu"]`, so a literal `auto` would make argparse exit rc=2. M4 resolves it before argv on **both** entry orderings:

- run_predict builds `--device` from `device = self.\_effective\_device()` (`native\_wasb\_runner.py:221-224`), never `self.cfg.device`.
- Healthcheck-first: cache set at `:176`, reused with no 2nd probe. Healthcheck-skipped: `\_effective\_device`? `\_cuda\_is\_available` lazily probes (`:122-126`). Either way argv gets cuda/cpu. Pinned by `test\_run\_predict\_auto\_threads\_cpu\_when\_no\_gpu` (`assert "auto" not in argv`) and `test\_run\_predict\_auto\_threads\_cuda\_when\_gpu`. The two metadata sinks (`fusion\_hybrid.py:87`, `wasb\_hybrid.py:46`) only record `cfg.device` for reporting ? they never build argv.

3. PROBE CACHING & CORRECTNESS ? PASS

- Coherent: healthcheck sets `self.\_cuda\_available = "cuda True" in out` (`:176`) for every device once torch imports, from a real probe; `\_cuda\_is\_available` (`:122-126`) reuses it (None?probe-once memoization). GPU availability can't change within a process, so no staleness. Production always runs healthcheck before run_predict on the same instance (`trajectory\_hybrid.py:228,239`), so the lazy path is the test/direct-API case only.
- `\_probe\_cuda` (`:112-120`) treats python-missing/timeout (caught `FileNotFoundError`/`TimeoutExpired` ? `False`), import-error/rc!=0 (`res.returncode == 0` and ... ? `False`) all as "no cuda" ? auto?cpu, never a crash. Pinned by `test\_healthcheck\_caches\_probe\_for\_run\_predict` (asserts `spy.call\_count` stays 1) and `test\_run\_predict\_explicit\_cuda\_does\_not\_probe` (`r.\_cuda\_available` is None).

4. EDGE CASES ? PASS

- `auto` + torch import fails (rc!=0) in healthcheck: returns False early at `:173-174`, `\_cuda\_available` stays None, production aborts. If healthcheck is skipped, the lazy `\_probe\_cuda` subprocess also fails ? resolves to `cpu`, then the inference subprocess fails and run_predict returns `None` gracefully (`:238-240`). No crash.
- `mps`: passthrough, no probe, no warning (`device.py` mps branch; `test\_explicit\_cpu\_and\_mps\_pass\_through`).
- Warning fires only on real fallback: `:182` if `ctx.fell\_back` (`device.py:43-46`) is True only for auto-without-cuda. Pinned by `test\_healthcheck\_auto\_falls\_back\_to\_cpu\_with\_warning`.

5. TEST GAPS / OS ? PASS

All five contracts are pinned by a failing-on-regression test: explicit-cuda-strict, auto?cpu, auto?cuda, no-probe-for-explicit, cache-reuse (enumerated above). `DeviceContext` frozen-ness (`test\_is\_frozen` ? `FrozenInstanceError`) and all three properties (`effective`/`fell\_back`/`cuda\_required\_but\_missing`) covered in `tests/test\_device\_context.py`. No OS-dependent assertions ? all subprocesses are mocked and Windows paths are normalized (`run\_predict :198`); the suite passed here on win32.

Nits (non-blocking, optional)

- **nit** ? Unknown device string (e.g. `"gpu"`): passes healthcheck (not "cuda", not "auto" ? no gate) and only fails later at `wasb\_infer.py` argparse (rc=2 ? run_predict returns None). This is pre-existing behavior M4 didn't regress, but since `DeviceContext` is now the device-policy seam, a future hardening could validate `requested` ? {auto,cuda,mps,cpu} in healthcheck for a clearer early error. Not required for M4.
- **nit** ? `\_probe\_cuda` (`native\_wasb\_runner.py:115`) and healthcheck (`:165`) use two slightly different print formats, but both yield the `"cuda True"` substring the parsers key on (verified). Harmless duplication; a shared probe-code constant would remove the chance of them drifting apart.

No edits made (review only).