

claude-windows / 7d6bbb83-42c0-459b-9e23-e14e96d4cb2d

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d.jsonl`  
- SHA-256: `f5b6e6d08177f7fabc21e0de1efd0e926c21bc20a36a4687c3217d44eac9d134`  
- Source modified: `2026-06-22T09:30:07+00:00`  
- Imported at: `2026-07-05T16:48:18+00:00`  
- project: `C--Users-avidu-Projects-badminton-highlight-indexer`  
- session_id: `7d6bbb83-42c0-459b-9e23-e14e96d4cb2d`
```

Transcript

1. user (2026-06-13T19:17:18.790Z)

Produce a DESIGN DOC at `docs/ANALYZERS\_AND\_RECONCILERS.md` (docs-only, NO implementation code; owner-gated; a peer to `docs/EXPERIMENT\_HARNESS.md` and `docs/MULTI\_SIGNAL\_FUSION\_PLAN.md`) for two extensibility seams in this badminton (? racket-sports) rally-detection pipeline.

FIRST: `git pull --ff-only` on master, then read for context ? `CLAUDE.md` (workflow conventions: branch off master, ONE PR per item, don't merge without owner), `docs/MULTI\_SIGNAL\_FUSION\_PLAN.md` (§2 the "RallyCue = swappable producer" pattern; §3.2 the scale/translation-invariant feature discipline), `docs/EXPERIMENT\_HARNESS.md` + `backend/eval/experiment.py` (the no-regression A/B + LOVO-honest measurement discipline: Variant, compare, compare_unlabeled), `docs/QUALITY\_ITERATIONS.md`, `docs/HOW\_RALLY\_DETECTION\_WORKS.md`, the fusion seam `backend/pipeline/segmenters/fusion\_hybrid.py` + `trajectory\_hybrid.py` (the `\_enrich\_candidate\_stats` and `\_select\_for\_handoff` hooks) + `backend/pipeline/detectors/fusion\_features.py` (pure-feature discipline) + `backend/pipeline/detectors/person\_detector.py::detections\_per\_frame`, and the events substrate in `backend/infrastructure/database.py` (`add\_event`, `add\_candidate\_segment`/`query\_candidate\_segments`, the `candidate\_segments.stats` JSON + the `events` table). (If `fusion\_hybrid.py` isn't on master yet it's the in-flight fusion-P2 PR ? design against `MULTI\_SIGNAL\_FUSION\_PLAN.md`; the seam is the per-window enrichment hook.)

THE OVERARCHING PRINCIPLE (owner directive ? make it the SPINE of the doc): ****NO special-casing in the decision path.**** Every analyzer/reconciler emits a SIGNAL carrying a confidence / presence; the SCORING MODEL (the learned classifier today, a richer scorer later) consumes those signals and **learns to weight them** ? we do NOT add hand-coded if/else branches into the decision. Analyzers propose; the scorer arbitrates; the experiment harness + golden set decide what gets promoted. This is what lets us pile on many clever signals (net-hit fault, let-cord, double-bounce, shuttle-out) without the pipeline becoming a pile of special cases.

DESIGN TWO SEAMS:

1) INJECTABLE WINDOW ANALYZERS (forward). A registry of pluggable `WindowAnalyzer`s (mirror `segmenter\_registry`) that run over a candidate window's PER-FRAME-ALIGNED signals ? shuttle (x, y, velocity, visibility) per frame + person boxes per frame + net line + window bounds, ALL already produced (WASB trajectory + `PersonDetector.detections\_per\_frame` + `rally\_gate`'s net estimate). Each emits an `AnalyzerResult` with up to three roles: (a) features (with confidence) merged into `candidate\_segments.stats` ? consumed by the scoring model + the experiment harness as columns; (b) events into the `events` table; (c) OPTIONAL deterministic boundary hints. Spec a concrete worked example: a SERVICE-FAULT analyzer ? shuttle enters the net region and its velocity halts (comes to rest) ? high-probability net-hit service fault ? emits a confidence-scored `service\_fault` signal + a `net\_hit` event + a rally-END boundary hint. The injection point is the existing FusionSegmenter `\_enrich\_candidate\_stats` hook (generalize "compute fixed features" ? "run each enabled analyzer"). Flag-gated, default OFF (via `cue\_flags`).

2) RECONCILIATION PASS (backward) ? a "linter for rally decisions". Given predicted segments + the persisted per-frame/per-window signals, detect where the DECISION contradicts the EVIDENCE

and propose corrections. Worked rules: shuttle-inactive-but-still-in-active-rally ? truncate to the last-active frame (the misfire fix the owner described); no-net-crossing rally ? drop/flag; over-merged (long dead-time gap inside one segment) ? split; boundary slack ? tighten to serve-contact. ****Report-first, apply-optional**** ? output an issues report (per segment: kind, evidence, before?after, reason) AND a corrected set; apply under a flag or surface for human review; NEVER silently rewrite (auditability + the repo's honesty/attribution working agreement). It must be expressible as an experiment-harness `Variant` transformation so its lift is A/B-measurable on the 6 golden videos BEFORE it is ever promoted to auto-apply. Idempotent; fixed rule precedence.

THE DOC SHOULD COVER: a unifying pipeline diagram (detectors ? analyzers ? decision/scoring ? reconcilers ? scoring/events/highlights); the `WindowAnalyzer` / `AnalyzerResult` / `ConsistencyRule` / `Issue` contracts (signatures, JSON-serializable, registry pattern); HOW signals flow into the scoring model as confidence-weighted features (NOT branches) ? tie explicitly to the learned classifier + harness; persistence (stats / events / a corrections record); the no-regression + harness-measured + flag-gated-default-OFF + rollback=config-flip discipline; how it composes with fusion P2 (analyzers plug into its enrichment hook, reconcilers consume its persisted signals); trade-offs/risks (deterministic-vs-learned signals, ordering/idempotence, signal-not-verdict, report-first auditability, small-N feature discipline); and a phased, owner-gated PR plan (each = ONE PR off master, default OFF, promoted only on measured LOVO lift). Cross-link it from `docs/README.md` and reference ADR-007 (modular rally layer). Keep it the size/shape of `EXPERIMENT\_HARNESS.md`.

Deliverable = the design doc on a branch + a docs-only PR (do NOT merge ? owner reviews). This is a DESIGN task: do not write pipeline implementation code.

2. assistant (2026-06-13T19:17:22.127Z)

I'll start by pulling latest on master, then reading the context files the task points me to.

3. user (2026-06-13T19:17:26.964Z)

Already up to date.

4. user (2026-06-13T19:17:27.013Z)

```
1      # Multi-signal rally fusion ? add a human (person) cue alongside the shuttle
2
3      **Status:** **P1 DONE** (2026-06-13) ? the person cue is extracted to
4      `backend/pipeline/detectors/person_detector.py` (`PersonDetector`), consumed by
5      `yolo_hybrid` with no
6      behavior change. **P2?P4 remain owner-gated.** Phases below are sequenced but each is a
7      separate PR.
8      **Owner decisions captured** in [S9](#9-open-decisions-for-the-owner).
9
10     > **Companion docs:** the 2-layer mental model is
11     [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md);
12     > the modular-fusion architecture this builds on is **ADR-007**
13     ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md));
14     > the permissive-dependency rule is **ADR-002 / ADR-005**; the audio cue (the #1 next
15     cue, already
16     scaffolded) is [archives/research/AUDIO_RALLY_DETECTION_PLAN.md](archives/research/AUD
17     IO_RALLY_DETECTION_PLAN.md).
18     > The vendor labels that gate the learned path come from
19     [GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md)
20     > and the labeling spec in [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md).
21     > **How each cue here is added/tested without production regression** ? A/B, shadow
22     eval, and the
23     > promote-only-on-lift discipline ? is [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)
24     (every P1?P4 cue
25     > below lands flag-gated, default OFF, promoted only on measured LOVO lift).
26
27     ---
28
29     ---
```

```

20     ## 1. Context ? why this, why now
21
22     Rally detection today runs one cue at a time. Each segmenter independently turns a
single signal
23     into rally windows:
24
25     | Segmenter | Signal | File |
26     |---|---|---|
27     | `wasb_hybrid` / `tracknet_hybrid` | shuttle velocity (frozen WASB/TrackNet trajectory)
| `backend/pipeline/detectors/tracknet_runner.py` |
28     | `yolo_hybrid` (the current `default_segmenter`) | person velocity (torchvision +
ByteTrack) | `backend/pipeline/segmenters/yolo_hybrid.py` |
29     | `motion` | raw pixel motion | `backend/pipeline/segmenters/motion.py` |
30     | `gemini` | a cloud VLM watching the whole clip | `backend/providers/gemini.py` |
31
32     There is no fusion. The owner's directive: make the detector use the shuttle and
the humans
33     together ? and be ready to fold in more signals (audio, later pose) ? delivered
34     MIT / Apache-2.0 / BSD-clean so we can ramp commercially without a licensing
rethink.
35
36     This is not greenfield ? it operationalizes work the repo already blessed:
37
38     - HOW_RALLY_DETECTION_WORKS.md open question: Fusion: combine shuttle-motion
(WASB) with
39     player-motion for static cams."
40     - ADR-007 already designed a modular "rally layer" that fuses cues outside the
frozen WASB
41     model, with audio as the #1 next cue and person/pose parked as #2. This plan picks
up the parked
42     person cue and generalizes the fusion seam so future cues are "register a producer,"
not "re-touch
43     the segmenter."
44     - backend/pipeline/detectors/rally_gate.py already exists ? the net-crossing gate,
explicitly
45     labelled in its own docstring as "Gate A in the over-fire fix spectrum (**B = player-
stance state
46     machine, future**)." The person cue in this plan is that named-but-unbuilt Gate
B.
47
48     ### The owner's "held-shuttle" false positive ? what's actually going on
49
50     The owner observed a "rally" that was just a player holding / flicking the shuttle in
hand between
51     points. This is the exact failure `rally_gate.py` was written to kill:
52
53     > WASB over-detects: when a player flicks/tosses the shuttle back to the opponent for
service after
54     > losing a point, that single trajectory looks like a rally. The discriminator is
structural: a real
55     > rally has the shuttle crossing the net MULTIPLE times; a single service feed crosses
~once."
56
57     The structural fix already exists and is validated offline ? but it is NOT wired into
the production
58     runtime path. backend/pipeline/detectors/rally_gate.py is called only from the
eval drivers
59     (`distill_local.py`, `eval_partial_labels.py`, `export_wasb_segments.py`,
`calibrate_local.py`); there
60     is no `min_crossings` knob in `IndexingConfig` and the live segmenters don't apply
it. So the
61     single cheapest, highest-confidence win in this whole plan is P2's first step: fold
the existing
62     Gate A into production. Everything else (the person cue / Gate B, the learned fusion)
is robustness

```

```

63     on top of that.
64
65     ### Why humans help ? and the honest risk
66
67     The beachhead videos are **static tripod** (Bangalore academies,
[PMF_MARKET_RESEARCH.md](PMF_MARKET_RESEARCH.md)),
68     so player presence/motion is a usable cue. ADR-001's warning was specifically about
**broadcast pan**
69     cameras, where player-motion fooled the old heuristic ? *that failure mode does not
apply to a fixed
70     court cam.* But we keep **the shuttle as the camera-invariant primary** so the general
bet stays robust
71     on other camera types. Humans then add:
72
73     - **Precision** ? reject shuttle-on-floor / warm-up knock-ups / held-shuttle feeds when
the court
74     isn't occupied by ?2 players in a play posture.
75     - **Recall** (via the learned classifier) ? features that help bridge shuttle
occlusion/blur gaps.
76
77     ---
78
79     ## 2. Core architecture ? one rally-fusion layer, cues are swappable producers
80
81     Formalize ADR-007's diagram. Every cue is a **black box that emits a time-aligned
signal** (per-frame
82     activity and/or per-window features); **fusion happens in *our* layer, never inside a
frozen model.**
83
84     ```
85     frames ?? WASB (frozen, MIT)      ?? shuttle trajectory ???
86     frames ?? person detector (ours)  ?? tracks + boxes ??????
87     audio  ?? hit detector (ours)     ?? hit events ???????????? RALLY LAYER (ours) ??
rallies
88     [frames ?? pose/stance (ours; parked #3) ?? stance ???????? feature + decision fusion
89     ```
90
91     - **Shuttle = primary, camera-invariant.** Windows still **seed** from
92     `TrackNetRunner.trajectory_to_action_windows()` (`tracknet_runner.py:234`). Other cues
*enrich and
93     adjudicate* shuttle-seeded windows; they do **not** replace the shuttle seed. Person
runs only over
94     **shuttle-seeded candidate windows + padding**, not the whole video ? which bounds the
added cost.
95     - **A small `RallyCue` contract.** Each producer yields `(frame_idx ? features)` + per-
window stats.
96     Adding a cue = register a producer; the existing `SegmenterRegistry`
(`segmenters/base.py:35`) is the
97     template for the pattern. Keep the single-cue segmenters registered as fallbacks.
98
99     ---
100
101     ## 3. The two fusion modes (both ? per the owner decision)
102
103     ### 3.1 Decision-level gate (safe baseline, ships first, no labels needed)
104
105     ```
106     rally_active = shuttle_motion ? (?? persons in court zone ? recent person track)
107     ```
108
109     The `? recent-track` clause tolerates 1?2 frame person-detector misses so the AND-gate
can't cause a
110     recall cliff. This composes with the **existing net-crossing Gate A**:
111
112     ```

```

```

113     keep window ? net_crossings ? 2           (Gate A ? EXISTS in rally_gate.py, eval-
only today)
114         ? (?? in-court players ? recent track)   (Gate B ? the person cue, NEW)
115     ...
116
117     Gate A kills the service-feed / held-shuttle case structurally; Gate B adds court-
occupancy evidence
118     for the cases a single trajectory can still fool (e.g. a knock-up that does cross the
net). No training
119     data required ? this is the label-free fusion that ships before any golden labels exist.
120
121     ### 3.2 Feature-level (primary, the learned path)
122
123     The cues emit a **small set of aggregated, scale/translation-invariant per-window
features ? NOT raw
124     coordinates.** Raw pixel coordinates would memorize *this* court/camera and fail to
generalize off a
125     still-small labeled set; counts / ratios / zone-membership / two-sidedness transfer.
126
127     **What already feeds the classifier** (`distill_local.py:40`, `FEATURE_NAMES`):
128     `duration, peak_velocity, mean_velocity, active_ratio, net_crossings` ? note
**`net_crossings` is
129     already a feature** (computed via `rally_gate.gate_windows`).
130
131     **Person features to add (~4, deliberately few):**
132
133     | Feature | Meaning | Rally signal |
134     |---|---|---|
135     | `players_in_court` | median count of persistent in-court tracks | ? 2 (singles) / 4
(doubles); 0?1 = dead time |
136     | `two_sided_motion` | are moving players on **both** net-halves? | rally = two-sided;
one person feeding = one-sided |
137     | `mean_player_motion` | aggregate normalized player velocity | rally = sustained;
resting/standing = low |
138     | `in_court_ratio` | fraction of frames with ?2 in-court players | track stability vs
flicker |
139     | `shuttle_player_colocation` | fraction of frames the shuttle sits inside/adjacent to a
person box (hand region) | **held = high; in-flight = mostly free space** ? the direct held-
shuttle discriminator |
140
141     These join the shuttle features in the candidate `stats` JSON
142     (`database.py::add_candidate_segment`) and the feature vector of the **existing
distilled classifier ?
143     the tiny numpy-only L2 logistic regression in `backend/eval/classifier.py` (ADR-004),
NOT a new net.**
144     It scores each candidate window `P(real rally)` and replaces the Gemini call at runtime.
145
146     **What the model learns:** the **weights** on those few features ? e.g. down-weight a
shuttle-motion
147     window with 0?1 in-court players / one-sided motion / high shuttle-colocation (warm-up,
floor shuttle,
148     held-shuttle feed); up-weight ?2 players, two-sided, sustained motion, shuttle in free
flight.
149
150     **Small-N discipline (non-negotiable):**
151     - Keep to **~4 person features** ? more features on a still-small labeled set overfit;
that's exactly
152     why the model stays a logistic regression, not a net.
153     - Use the **leave-one-video-out (LOVO)** protocol that `distill_local.py` **already
implements**: train
154     on the other videos' candidates, predict the held-out *video*, score vs its labels ?
never a
155     held-out *window* from the same video (windows in one video are correlated; that
flatters the score).
156     - Training labels today are **Gemini silver labels** (distilled offline, ADR-004/007 ?

```

Gemini is a

```

157     teacher, never a runtime dependency, never a training-data *source* for owned
weights); **golden
158     labels are the honest measuring stick**, and as the Roora golden corpus grows they can
also become
159     training labels. Until the labeled set is larger, treat the learned weights as
160     **calibration / direction-check, not "a fully general model"*** ? the decision-level
gate (§3.1) is
161     the shippable, label-free fusion in the meantime.
162
163     > *(Gated on golden labels existing ? the same blocker ADR-007 flags for audio. See
164     > [GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md).)*
165
166     ---
167
168     ## 4. Shuttle-state: in-flight vs in-hand (the misclassification, mapped to real code)
169
170     The "shuttle moved" seed conflates a shuttle **in flight** with a shuttle **carried in
hand**. The
171     discriminators, and where each lives:
172
173     | Signal | In-flight | Held / static | Status |
174     |---|---|---|---|
175     | `net_crossings` | ? 1 (usually several) | ~0 (held) / ~1 (single feed) | **EXISTS**
(`rally_gate.py`); eval-only ? wire to prod |
176     | `shuttle_player_colocation` | low (free space) | high (near a hand) | **NEW** (needs
the person cue) |
177     | `direction_reversals` | several (racket contacts) | ~0 | derivable from trajectory;
partial overlap with crossings |
178     | `peak shuttle speed` | high | walking speed | derivable from existing velocity |
179
180     So the in-flight-vs-held fix is **mostly Gate A (already built) + the new colocation
feature from the
181     person cue**. **No new label column is needed** ? the golden labels already start a
rally at the
182     *serve-contact frame* and treat all pre-serve / held-shuttle time as dead time (see
183     [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md)), so held-shuttle moments
are negatives
184     by construction and are the ground truth that proves the detector wrong.
185
186     ---
187
188     ## 5. Adversarial mitigations (integrating safely)
189
190     - **Court-region mask** ? polygon around the court; drop persons outside it (spectators,
coaches,
191     adjacent courts). Config-supplied per camera; ~1 ms. Source: a manual per-camera
polygon first;
192     the **4 court corners + net-line** that Roora labels once per video (see the labeling
doc) build
193     this mask **and** the net split that `two_sided_motion` / `net_crossings` need.
194     - **Player-count constraint** ? expect **2 (singles) / 4 (doubles)** persistent in-court
tracks;
195     suppress transient/extra detections. The `singles_or_doubles` manifest flag sets the
expectation.
196     - **Temporal association** ? reuse `supervision` **ByteTrack** (already in
`yolo_hybrid`); reject
197     singleton detections; smooth across frames.
198     - **Multi-court rejection** ? proximity to the primary court region (homography
optional, later).
199     - **Domain shift (COCO frontal ? side/overhead court)** ? start **box-only** (robust);
keep the
200     detector swappable. Fine-tuning a person detector is **parked** ? we have no in-house
person labels,
201     and the **paid-LLM/training guardrail + minors-exclusion** apply to any future person-

```

```

training data.
202     - **Fusion-degradation guard** ? shuttle stays primary; person is gate + feature, never
the sole
203         trigger; the single-cue segmenters remain registered fallbacks.
204
205     ---
206
207     ## 6. Box vs pose
208
209     **Box-first.** Occupancy + count + two-sided motion + shuttle-colocation is enough for
both the gate
210     and the features, and box detection is the cheap, robust win. **2D pose is parked behind
the same
211     `RallyCue` interface ? pose (MediaPipe Pose / RTMPose, both Apache-2.0) buys serve-
ready stance
212     gating and shot cues later, at the cost of a second model. This matches ADR-007 parking
pose as the
213     #2/#3 cue.
214
215     ---
216
217     ## 7. Performance
218
219     The person detector is the **existing in-process torchvision model** run at **frame-skip
3?5**
220     (`yolo_frame_skip` default 3) with track interpolation between sampled frames. The
shuttle keeps
221     running in its separate WSL process ? the two are **separate processes**, so a "shared
backbone" isn't
222     free; budget realistically as the existing `yolo_hybrid` cost minus skip savings. Person
runs **only
223     over shuttle-seeded windows + `ai_handoff_padding`, not the whole video, which bounds
the cost.
224
225     ---
226
227     ## 8. Licensing (the guardrail table ? all green)
228
229     Every dependency stays **MIT / Apache-2.0 / BSD** (ADR-002 / ADR-005; `ultralitics` was
dropped for
230     AGPL-3.0).
231
232     | Cue / model | Code | Weights / data | Verdict |
233     |---|---|---|---|
234     | Shuttle: WASB-SBDT, TrackNetV3 | MIT | frozen badminton weights | ? (ADR-001/002) |
235     | Person (default): torchvision Faster R-CNN / RetinaNet / FCOS | BSD/MIT | COCO-
pretrained | ? already in repo (`_DETECTOR_FACTORIES`) |
236     | Person (escalation): **YOLOX**, **RT-DETR / RT-DETRv2** | Apache-2.0 | COCO /
Objects365 | ? |
237     | Pose (parked): **MediaPipe Pose**, **RTMPose / MMPose** | Apache-2.0 | COCO-structure
| ? |
238     | Tracking: `supervision` ByteTrack | MIT | n/a | ? already in repo |
239     | **Banned** | **Ultralytics YOLOv5/8/11 (AGPL-3.0)** . **MonoTrack (Adobe, non-comm)**
. **YOLO-NAS (Deci, non-comm)** | ? | ? |
240
241     **COCO caveat to record:** COCO *annotations* are **CC BY 4.0** (commercial OK *with
attribution*);
242     COCO *images* are Flickr-ToS (per-image). Pretrained-**weight** use is clean; record the
attribution.
243
244     ---
245
246     ## 9. Open decisions (for the owner)
247
248     1. **Court-region mask source** ? manual per-camera polygon (simple, ship first) vs auto

```

```

court
249     homography/line detection (later). **Recommend manual first**, fed by Roora's court-
corner labels.
250     2. **Person detector default** ? keep torchvision (zero new deps, BSD/MIT) vs adopt
YOLOX/RT-DETR
251     (Apache, faster) now. **Recommend torchvision first**, escalate only if eval shows a
speed/accuracy gap.
252     3. **Pose now or parked** ? **recommend parked** behind the interface (box-first).
253     4. **Golden labels** ? feature-level fusion (P3) is **blocked on golden rally labels**
the source is
254     the VLC `rally-annotator` + the Roora vendor pilot
([GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md)).
255
256     ---
257
258     ## 10. Phased execution (owner-gated; ONE PR per slice, off `master`)
259
260     - **P0 ? this design doc** +
[ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md), indexed in
261     `docs/README.md`, cross-linked from `HOW_RALLY_DETECTION_WORKS.md` and ADR-007. *(this
deliverable;
262     docs-only, no code.)*
263     - **P1 ? extract the shared person cue:** ? **DONE (2026-06-13).** **Lifted the
torchvision detector +
264     ByteTrack + velocity out of `yolo_hybrid` into
`backend/pipeline/detectors/person_detector.py`
265     (`PersonDetector`); `yolo_hybrid` now orchestrates only and delegates to `self.person`
(no behavior
266     change ? regression-safe, proven by the unchanged process_video output tests). New
267     `tests/test_person_detector.py`; suite +1 green; person_detector is type-clean (left
mypy quarantine
268     for the moved code). The fusion segmenter (P2) reuses this same producer over shuttle-
seeded windows.
269     - **P2 ? wire Gate A into production, then add the court-mask + player-count gate (Gate
B):** add a
270     `min_crossings` knob + court-region polygon + count constraint to `IndexingConfig`; a
271     `FusionSegmenter` (or extend the hybrid segmenter) runs the person cue over shuttle-
seeded windows
272     and applies `Gate A ? Gate B`. Persist person features into the candidate `stats`.
**Wiring Gate A
273     is the immediate false-positive win.**
274     - **P3 ? feature-level (learned) fusion:** add the person features (incl.
`shuttle_player_colocation`)
275     to `distill_local` / `classifier`; recalibrate IoU on the golden set; report **LOVO
lift on held-out
276     videos**. *(Gated on golden labels.)*
277     - **P4 ? additional cues:** wire **audio** (ADR-007's #1, already scaffolded in
278     `backend/pipeline/audio/hit_detector.py`) and later **pose**, through the same
`RallyCue` interface.
279     Each is add-a-producer, not a refactor.
280
281     ---
282
283     ## 11. Verification
284
285     - **Offline, runnable in this dev container (no GPU/torch/cv2):** unit tests for the
fusion gate logic
286     and `RallyCue` aggregation (court-mask filter, player-count constraint, `? recent-
track` tolerance)
287     and the shuttle-state signals (net-crossing count ? extends the existing `rally_gate`
tests ? and
288     person-colocation) with **synthetic trajectories + person boxes**. A key-license
assertion test
289     (no `ultralytics` import; detector factories stay torchvision/Apache).
290     - **Machine-side (owner, GPU/WSL ? the full pipeline can't run in this container):** run

```



```

the
291     `FusionSegmenter` on a real static-tripod clip; confirm **held-shuttle / warm-up /
spectator false
292     positives drop** vs shuttle-only (this should reproduce the owner's observed bug and
show it fixed ?
293     start by confirming Gate-A-in-production alone fixes it); **LOVO IoU lift on the
golden set** for the
294     learned path.
295     - **Docs (P0):** links resolve; both docs indexed in `docs/README.md`; this doc cross-
linked from
296     `HOW_RALLY_DETECTION_WORKS.md` (the fusion open-question) and ADR-007.
297

```

5. user (2026-06-13T19:17:27.027Z)

```

1     # Experiment harness ? A/B + shadow eval for rally-detection quality (no-regression by
design)
2
3     **Status:** **P1 IMPLEMENTED** (`backend/eval/experiment.py` +
`tests/test_experiment.py`, shipped
4     2026-06-13) ? the `Variant` layer, `compare()` (labeled A/B, LOVO-honest),
`compare_unlabeled()`
5     (SxS on NEW unlabeled footage), and the leaderboard. **P2?P4 remain owner-gated.**
Phases below are
6     sequenced but each is a separate PR.
7
8     > **Companion docs:** the cues this experiments *on* are in
9     > [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md); the scoring mechanics are
10    > [EVALUATION.md](EVALUATION.md); the no-regression gate + golden labels +
`baseline.json` come from
11    > [GOLDEN_SET_VENDORING.md](GOLDEN_SET_VENDORING.md); the modular-cue architecture is
**ADR-007**
12    > ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the distilled
classifier is
13    > **ADR-004**. The 2-layer mental model is
[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md).
14
15    ---
16
17    ## 1. Context ? why this, why now
18
19    PR #112 (merged) added the multi-signal fusion design ? shuttle + person + audio cues,
all **default
20    OFF**. The owner's directive for *how* we evolve quality from here:
21
22    > **"Evolve the codebase so we can keep finding new ways to improve quality and
experiment with them ?
23    > or add them as a signal ? without suffering any regressions if the idea doesn't work
out. Rolling
24    > back the code would be too dangerous."**
25
26    So **experimentation and deployment must be decoupled.** The contract:
27
28    1. Merge experimental code freely ? it's **gated OFF**, so it cannot change production
behavior.
29    2. Test it **offline against golden labels** (most experiments never touch the served
path at all).
30    3. Promote a variant to default **only on measured lift**, through a CI eval gate.
31    4. **"Rollback" = flip a config flag, never `git revert`.** The proven path is a pinned
variant that is
32    always present.
33
34    **The pleasant surprise (Explore audit, 2026-06-13):** ~80% of this harness already
exists and is
35    simply not composed *as a system*. The missing piece is a thin **variant/experiment

```

```

layer + the
36     discipline**, not a platform. §4 is the exact code surface so the next session does not
re-discover it.
37
38     ---
39
40     ## 2. The thesis ? separate expensive DETECTION from cheap DECISION
41
42     The expensive, GPU/WSL-bound, risky part is **detection**: "what signals exist per
frame" (shuttle
43     WASB trajectory, person tracks/boxes, audio hits). The cheap, swappable, pure-Python
part is the
44     **decision**: "given those signals, where are the rallies" (windowing + gates + the
classifier).
45
46     > **Run detection once per video, persist all per-cue signals + per-candidate features,
then re-score
47     > any number of variants OFFLINE ? deterministically, with no GPU and zero production
risk.**
48
49     This is not aspirational ? `distill_local.py` and `calibrate_local.py` already do
exactly this on
50     stored WASB trajectories. The candidate-segments table (`stats` JSON) is the persistence
substrate.
51     ? Most experiments are pure re-scoring of stored data; they never enter the served
pipeline.
52
53     ---
54
55     ## 3. What already exists (the audit ? do NOT rebuild)
56
57     Exact public surface, with `file:line`. **This table is the anti-duplicate-work
artifact** ? start here.
58
59     ### 3.1 Scoring (variant-agnostic ? grades any `List[Interval]`)
60     `backend/eval/metrics.py`
61     - `interval_iou(a, b) -> float` (:15) ? temporal IoU of two `(start,end)` intervals.
62     - `match_intervals(preds, gts, iou_threshold=0.5) -> MatchResult` (:48) ? greedy 1-to-1
IoU matching.
63     - `score(preds, gts, iou_threshold=0.5, match_result=None) -> Score` (:104) ? P/R/F1,
mean IoU, boundary errors; `Score.as_dict()`
64     - `pr_curve(preds, gts, thresholds=None) -> List[Score]` (:138).
65
66     ### 3.2 A/B + cross-validation primitives (the comparison engine already exists)
67     `backend/eval/calibration.py`
68     - `improvement_report(predict_fn, baseline_config, tuned_config, gts, iou_threshold=0.5)
-> dict` (:148) ? **before/after metrics + %?. This IS the A/B delta.**
69     - `leave_one_video_out(videos, configs, metric="f1", iou_threshold=0.5) -> dict` (:113)
? **LOVO: pick on other videos, score on held-out video (honest cross-video).**
70     - `cross_validate(predict_fn, configs, gts, k=5, metric="recall", iou_threshold=0.5) ->
dict` (:72) ? within-video k-fold overfit guard.
71     - `select_best(predict_fn, configs, gts, metric="f1", iou_threshold=0.5) -> (Config,
float)` (:61).
72     - `evaluate(preds, gts, iou_threshold=0.5) -> Score` (:41); `kfold_indices(n, k)` (:46).
73     - LOVO video dict shape: `{"name": str, "predict_fn": PredictFn, "gts": [Interval]}`.
74
75     ### 3.3 No-regression gate + baselines (the promotion gate already exists)
76     `backend/eval/regression.py`
77     - `regress(db, video_id, ground_truth, stage="final", iou_threshold=0.5, detector=None,
tolerance=DEFAULT_TOLERANCE, baseline_path=DEFAULT_BASELINE_PATH) -> RegressionResult` (:104) ?
score stored preds vs GT, flag if P **or** R drops > tolerance.
78     - `regress_with_provider(...)` (:160) ? same, fetching GT via the annotation provider.
79     - `save_baseline(video_id, stage, precision, recall, f1, ..., gt_fingerprint=None)`
(:80) . `load_baselines(path) -> dict` (:68).
80     - `RegressionResult` (:43): `regressed`, `reasons`, `gt_hash`, baseline P/R,

```

```

`tolerance`, `has_baseline`; `as_dict()`.
81     - Default baseline file: `eval_baselines/regression_baseline.json` (:33).
82
83     `run_regression.py` ? CI-gatable CLI, `main(argv)` (:52). **Exit codes: 0 ok / 1
regression / 2 no-DB / 3 no-baseline.** Flags: `--video-id --tier --annotations
--stage{candidates,final} --detector --iou --tolerance --baseline --update-baseline --allow-
missing-baseline --json`.
84
85     ### 3.4 Pinnable model artifact
86     `backend/eval/classifier.py` ? `LogisticRegression` (pure numpy, L2):
`fit(X,y,feature_names=,class_weight="balanced")` (:40), `predict_proba` (:75),
`predict(threshold=0.5)` (:81), `to_dict`/`from_dict` (:86/:97), `save`/`load` (JSON)
(:107/:111). **A trained model = one JSON file ? a pinnable, hashable variant artifact.**
87
88     ### 3.5 Offline re-scoring substrate (persist once, re-score forever)
89     `backend/infrastructure/database.py`
90     - `add_candidate_segment(video_id, detector, start_time, end_time, chunk_index=None,
confidence=None, stats=None, detection_params=None) -> Optional[int]` (:240) ?
`stats`/`detection_params` stored as JSON; `candidate_segments` table cols incl. `stats`,
`detection_params`, `human_label`, `confirmed_segment_id` (schema :97).
91     - `query_candidate_segments(video_id, detector=None, only_unlabeled=False) ->
List[dict]` (:277) ? returns rows with deserialized `stats`.
92
93     ### 3.6 Label assignment + feature glue
94     `backend/eval/training.py`
95     - `label_candidates(candidate_intervals, gt_intervals, iou_threshold=0.5) -> List[int]`
(:50) ? y-labels by IoU match to golden (same matcher as the evaluator).
96     - `build_training_set(candidate_rows, gt_intervals, iou_threshold=0.5,
feature_fn=extract_yolo_features) -> (X, y, intervals)` (:64) ? pluggable `feature_fn`.
97
98     ### 3.7 Existing LOVO drivers to mirror (not duplicate)
99     - `backend/eval/distill_local.py` ? `run(specs, iou=0.5, threshold=0.5, model_out=None)`
(:98), CLI `python -m backend.eval.distill_local --manifest videos.json --model-out ...`;
`FEATURE_NAMES = [duration, peak_velocity, mean_velocity, active_ratio, net_crossings]` (:40);
already does LOVO over Gemini silver labels.
100    - `backend/eval/calibrate_local.py` ? `run(specs, iou=0.5, metric="f1")` (:77);
`LocalConfig = (velocity_thresh, merge_gap, min_window_duration, min_crossings)`; grid + LOVO.
101    - `backend/eval/calibrate_wasb.py` ? `run(...)` , `load_golden_gts()`;
`backend/eval/gt_loader.py::load_gt_intervals(csv_path, *, strict=True)` is THE canonical golden
reader.
102
103    ### 3.8 Registries + config knobs (variant selection)
104    - `backend/pipeline/segmenters/base.py` ? `SegmenterRegistry` (:35), singleton
`segmenter_registry` (:63); registered: `motion`, `gemini`, `yolo_hybrid`, `tracknet_hybrid`,
`wasb_hybrid`.
105    - `backend/annotations/base.py` ? `AnnotationProviderRegistry`, singleton
`annotation_registry`; providers `file`, `auto`.
106    - `backend/config/models.py::IndexingConfig` ? `default_segmenter` (:65, default
`"yolo_hybrid"`), `detector_model` (:72), `detector_score_threshold` (:73),
`yolo_velocity_threshold` (:74), `yolo_frame_skip` (:75), `min_segment_duration` (:68),
`dead_time_merge_gap` (:69), `motion_threshold` (:67). **No `min_crossings` knob yet** (see
fusion P2).
107
108    ### 3.9 Confirmed ABSENT (no duplication risk)
109    Explore found **no** existing experiment / variant / shadow / A-B *machinery* ? only a
stray "A/B test"
110    aspiration comment in `backend/pipeline/detectors/base.py` and an "experiment E1" note
in
111    `AUDIO_RALLY_DETECTION_PLAN.md`. The `regress()` + baseline path is the canonical
comparison mechanism.
112
113    ---
114
115    ## 4. The variant abstraction (the one thin thing to add)
116

```

```

117     A **`Variant` = a declarative recipe**, not a code branch:
118
119     ```
120     Variant = (segmenter_name, IndexingConfig overrides, enabled cues + per-cue params,
classifier model path/hash)
121     ```
122
123     JSON-serializable; selected via the existing `segmenter_registry` +
`IndexingConfig.default_segmenter`.
124     **Control** = a *pinned, frozen* variant (recipe + model hash) that is always registered
and is the
125     default. Adding a new cue = a new Variant differing by one flag ? the control recipe is
untouched.
126
127     ---
128
129     ## 5. Three experiment modes
130
131     ### 5.1 Offline A/B (primary, label-driven, zero production risk)
132     `compare(videos, variant_A, variant_B)` ? a thin driver composing **existing**
functions:
133     `query_candidate_segments()` (load persisted features) ? re-score each variant ?
`calibration.improvement_report()`
134     + `calibration.leave_one_video_out()` + `metrics.score()`, graded vs golden
(`gt_loader.load_gt_intervals`).
135     Reports per-video **and** LOVO-honest aggregate IoU/P/R/F1 deltas. **This is the A/B
test, and it needs
136     almost no new scoring code** ? only the glue.
137
138     ### 5.2 Shadow (for cues that touch detection, e.g. the person detector)
139     The new cue runs and **persists its signal** into the candidate
`stats`/`detection_params`, but the
140     *served* decision stays control. Offline you then ask "what would variant B have
scored?" via §5.1 ?
141     B never affects output until it wins. This is how a detection-touching change (not just
a re-scoring
142     change) earns its way in without risk.
143
144     ### 5.3 Promotion gate (no-regression, CI-enforced)
145     `regression.regress()` + `baseline.json` + `run_regression.py` exit codes. A variant
becomes the new
146     default **only if** it doesn't regress the golden set beyond `tolerance` (exit 1
blocks). Control is
147     pinned ? **rollback = flip `default_segmenter`/variant config, never a code revert.**
148
149     ---
150
151     ## 6. No-regression guarantees (the heart of the owner's ask)
152
153     - **New cues/signals default OFF** ? per-cue enable flag in `IndexingConfig`; merged
code can't change
154     behavior until explicitly enabled.
155     - **Control variant pinned** ? frozen recipe + classifier model hash, always registered
& default.
156     - **Promotion only via the eval gate** ? `run_regression.py` in CI (exit 1 blocks); no
silent default change.
157     - **Experiments are records** ? variant-spec hash ? per-video + LOVO scores + label-set
hash + timestamp,
158     persisted to a small JSON **experiment leaderboard** (generalizes `baseline.json`).
Answers "which
159     signals ever helped, on which video slices," and makes every result reproducible.
160     - **Decoupled events** ? "merge an experiment" and "change the default" are separate,
independently
161     revertible actions.
162

```

```

163     ---
164
165     ## 7. What to build later (gap list ? thin, additive, owner-gated)
166
167     1. ~~~`backend/eval/experiment.py`~~~ ? a `Variant` dataclass + `compare(videos, a, b)`
composing the $3
168         functions (no new scoring math).~~~ **DONE (2026-06-13).** Shipped `Variant` (+ pinned
`CONTROL`),
169         `compare` (labeled A/B with LOVO-honest learned-variant training, mirroring
`distill_local`),
170         `compare_unlabeled` (the **SxS-on-new-videos** mode ? agreed / A-only / B-only via
`match_intervals`,
171         no GT), `record_experiment` (? `eval_baselines/experiment_leaderboard.json`), and
`load_video_candidates`
172         (the `query_candidate_segments` bridge). JSON variant recipes first as recommended; a
`VariantRegistry`
173         only if the count grows. 19 tests incl. the no-regression invariant.
174     2. **Persist per-cue features** into candidate `stats` (person: `players_in_court`,
`two_sided_motion`,
175         `in_court_ratio`, `shuttle_player_colocation`; shuttle-state: `net_crossings`,
`direction_reversals`;
176         audio: hit times) so any future variant re-scores offline ? ties to
`MULTI_SIGNAL_FUSION_PLAN.md` $3.2/$4.
177     3. **Per-cue enable flags** in `IndexingConfig` (default `False`); `default_segmenter`
already exists.
178         Add the `min_crossings` knob (fusion P2) here too.
179     4. **CI promotion lane** using `run_regression.py` against the golden `baseline.json`;
write the
180         experiment-leaderboard JSON.
181
182     ---
183
184     ## 8. Phasing (owner-gated; ONE PR per slice, off `master`)
185
186     - **P0 ? this doc**, indexed in `docs/README.md`, cross-linked from
`MULTI_SIGNAL_FUSION_PLAN.md` and
187         `GOLDEN_SET_VENDORING.md`. *(this deliverable; docs-only, no code.)*
188     - **P1 ? `experiment.py` `compare()`** ? **DONE (2026-06-13):** thin wrapper over
`calibration` +
189         `metrics` + `query_candidate_segments`; pure-logic, fully tested in this dev container
(`compare`,
190         `compare_unlabeled`, `record_experiment`, the no-regression invariant).
`compare_unlabeled` adds the
191         SxS read on **new unlabeled footage** the owner asked for (run control + candidate on
a fresh video,
192         diff where they agree/diverge ? two reels via the existing `VideoCompiler` is the GPU-
side follow-on).
193     - **P2 ? per-cue enable flags + feature persistence** into candidate `stats` (default
OFF) ? couples
194         with fusion P1/P2.
195     - **P3 ? CI promotion gate**: `run_regression.py` lane on the golden `baseline.json`;
experiment-leaderboard JSON.
196
197     ---
198
199     ## 9. Scope guardrail
200
201     **Single-tenant / local-first** (per `PLATFORM_ARCHITECTURE.md`): a JSON experiment
registry + the
202     existing harness ? explicitly **NOT** MLflow / W&B / a service. Keep it the size of
`baseline.json`.
203
204     ---
205
206     ## 10. Verification

```

```

207
208 - **Offline, runnable in this dev container (no GPU/torch needed for the glue):** unit
test that
209 `compare()` is deterministic on synthetic candidate rows; a **no-regression invariant
test** asserting
210 "all new-cue flags OFF ? predictions byte-identical to control"; a CI-gate test
asserting
211 `run_regression.py` exit codes (0 ok / 1 regress / 3 no-baseline).
212 - **Machine-side (owner, GPU/WSL):** persist candidates once on a real clip, run
`compare(control,
213 fusion)` LOVO, confirm the reported lift matches a manual `metrics.score()` spot-
check.
214 - **Docs (P0):** links resolve; indexed in `docs/README.md`; cross-linked from
215 `MULTI_SIGNAL_FUSION_PLAN.md` and `GOLDEN_SET_VENDORING.md`.
216 </content>
217

```

6. assistant (2026-06-13T19:17:31.275Z)

Master is up to date. Now let me read the remaining context files in parallel.

7. user (2026-06-13T19:17:32.231Z)

```

1      # How rally detection works (the 2-layer mental model)
2
3      A short conceptual map of how we go from raw video to rally `[start, end]` segments ?
4      and, crucially, **which layer our annotation can improve.**
5
6      ```
7      video ??? frames ??? [DETECTION] ??? shuttle (x,y,visible) per frame
8              ?
9              (WASB CNN, per-frame trajectory
10             frozen weights) ?
11             ?
12             velocity signal (how fast the shuttle moves)
13             ?
14             [SEGMENTATION] ??? the 3 knobs we tune
15             ?
16             ?
17             rally windows [start, end] ??? IoU vs golden
labels
18     ```
19
20     ## Two layers ? keep them separate in your head
21
22     **1. DETECTION ? "where is the shuttle, this frame?"**
23     WASB is a CNN that, for each frame, outputs a heatmap ? the shuttle's `(x, y)` + a
24     visibility/confidence. Its weights were learned (by gradient descent, on shuttle-
coordinate
25     labels) and we use them **frozen**. We chose a *shuttle* detector on purpose: "shuttle
in
26     flight" is the signal for "a rally is happening," and it's **camera-agnostic** ? unlike
27     player-motion, which fooled the old heuristic on broadcast cameras that pan with the
players.
28     (Your videos are static tripod, so player-motion would also work ? but shuttle-motion is
the
29     robust general bet.)
30
31     **2. SEGMENTATION ? "when is a rally?"**
32     From the per-frame `(x, y)` we compute the shuttle's **velocity** (displacement between
frames).
33     Then three knobs turn that into rally intervals:
34     - **`velocity_thresh`** ? a frame is "active" only if the shuttle is moving faster than
this
35     (in flight, not lying on the floor / being picked up).

```

```

36     - merge_gap ? active frames within this many seconds get joined into one rally
(bridges
37     brief gaps: a shuttle momentarily occluded, or the beat between two hits).
38     - min_window_duration ? throw away windows shorter than this (blips / false
flickers).
39
40     Output: a list of [start, end] windows = the rallies.
41
42     ## Where the annotation fits (the key point)
43     Our Roora/golden labels are rally [start, end] times ? they say when rallies
are, not
44     where the shuttle is*. So they can't teach WASB's detector (that needs (x,y) labels).
They
45     can grade and tune the segmentation layer: we score predicted windows against
the golden
46     rallies with temporal IoU (how much they overlap), and pick the 3 knobs that
maximise it ?
47     calibrate on some rallies, measure on held-out ones so the number is honest.
48
49     ## Where sparks could land (open questions ? not commitments)
50     Things worth a shower-thought; any of these could beat 3 hand-tuned thresholds:
51     - Richer features than raw velocity: acceleration, trajectory continuity (a real
rally is a
52     smooth arc that survives occlusion), shuttle height/arc, or direction reversals
(each
53     reversal ? a shot ? ties to shots_count!).
54     - A tiny learned segmenter instead of 3 thresholds: feed the trajectory features to
a small
55     model trained on the golden rally labels (still the segmentation layer, just smarter).
56     - Hysteresis (two thresholds: enter-rally high, stay-in-rally low) ? robust to mid-
rally dips.
57     - Rhythm/periodicity: rallies have a hit cadence; dead time doesn't.
58     - Fusion: combine shuttle-motion (WASB) with player-motion for static cams. ?
designed in
59     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (one rally-fusion layer;
the person cue is
60     ADR-007's parked #2 and rally_gate.py's named-but-unbuilt "Gate B").
61     - Per-video calibration as a feature: "watch a few labeled rallies ? auto-tune to
this
62     court/lighting ? segment the rest sharply" ? exactly your earlier idea, at the right
layer.
63
64     The whole point of per-video calibration: a handful of your labeled rallies nudges
these knobs
65     to your court/camera, and we report the IoU lift on the rallies it never saw.
66

```

8. user (2026-06-13T19:17:32.713Z)

```

1     # Rally-detection quality iterations ? living log
2
3     Purpose. A running, reverse-chronological log of how we evolve rally-detection
quality: each
4     iteration's hypothesis, what shipped, the measured (or expected) effect, and the lesson.
This is the
5     live continuation of the historical trail in
6     [archives/quality-iterations/](archives/quality-iterations/README.md) ? per the
archive policy
7     (CLAUDE.md), terminal-state (DONE/REJECTED) work gets a dated subdir there; this
file stays at
8     docs/ top level and keeps growing, then its entries graduate to the archive when they
conclude.
9
10    Methodology (non-negotiable). Every quality change rides the experiment harness
11    ([EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)): new cues land flag-gated, default

```

```

OFF**; they're
12     graded **offline against golden labels** (per-video + LOVO-honest); promotion to default
happens **only
13     on measured lift** through the `run_regression.py` gate; **rollback = flip a config
flag, never
14     `git revert`.** Cues are designed in
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md); scoring
15     in [EVALUATION.md](EVALUATION.md); the 2-layer model in
16     [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md).
17
18     **Golden corpus = the 6 labelled matches** in `output/human_lovo_manifest.json` (owner-
confirmed
19     2026-06-13: use all 6 ? `adarsh_avi_singles, kushagra_singles, targetest_doubles,
gbaaddy,
20     testlarge_short, mahadevpura_singles`). The model trained on these is what we run
alongside the existing
21     detector on NEW footage for the side-by-side review.
22
23     ---
24
25     ## Baselines that drive strategy (IoU 0.5 vs human golden labels)
26
27     | Path | Mahadevpura (clean-ish) | Kushagra (multi-court) |
28     |---|---|---|
29     | Heuristic (velocity windowing) | 0.657 | 0.157 |
30     | Gen-0 owned head (LOVO, n=6) | 0.744 | 0.561 |
31     | Two-stage WASB+Gemini refine | 0.83 | ? |
32     | Gemini wrapper (`gemini-flash-latest`) | 0.93 | 0.66 |
33
34     **Reading:** clean footage is commoditized (serve via Gemini); **multi-court drops the
commodity wrapper
35     to 0.66 ? that is the owned model's ship target.** Baselines tracked in
`eval_baselines/`. The owner's
36     observed false positive ? a player **holding/flicking** the shuttle between points
counted as a rally ?
37     is the precision bug the fusion gates target.
38
39     ---
40
41     ## Current track (live) ? newest first
42
43     ### Multi-signal fusion P2: the FusionSegmenter *(in progress, 2026-06-14)*
44     **Hypothesis.** Adjudicating shuttle-seeded windows with the net-crossing **Gate A**
(`min_crossings?2`)
45     and a person-occupancy **Gate B** raises **precision** on the held-shuttle / warm-up /
one-sided-feed /
46     spectator false positives ? *without* hurting recall, because the shuttle stays the
primary,
47     camera-invariant seed.
48
49     **Scope ? and an honesty boundary from the archive (read this).** The
50     [2026-06-multivideo-lovo](archives/quality-iterations/2026-06-multivideo-
lovo/README.md) iteration
51     already smoke-tested **court-region detection as a fix for the *multi-court contrast*
confound** and
52     found it **low-headroom**: only 12?14% of dead-time shuttle detections fall *outside*
the foreground
53     court, so a perfect court gate lifts contrast to only ~1.6?2.2x (below the ?3.6x "works"
threshold) ? the
54     multi-court problem is **temporal**, and the **learned model**, not a court mask, is its
answer. So Gate B
55     here is **NOT** sold as the multi-court fix. It targets a **different, previously-
untested** failure mode
56     (held-shuttle/warm-up/one-sided/spectator FPs). Two distinct levers; don't conflate
them.

```



```

57
58     **Data support for Gate A.** The same archive shows real rallies on the multi-court
matches contain
59     **22 net-crossings 97% of the time (median 5)** while a single service feed crosses
~once ? `min_crossings?2`
60     is a principled, data-backed discriminator (it's the held-shuttle owner-bug killer).
61
62     **Shipping shape.** A NEW `fusion` segmenter, **default OFF** (`default_segmenter`
unchanged); Gate A
63     wired into the live trajectory path behind `indexing.min_crossings` (default 0 = OFF ?
wasb/tracknet
64     byte-identical today); the court mask is an **optional** config polygon (omit ? Gate B
is
65     player-count-only; supply ? masked). Person features (`players_in_court,
two_sided_motion,
66     mean_player_motion, in_court_ratio, shuttle_player_colocation`) persisted into candidate
`stats` for P3's
67     learned fusion. Reuses `PersonDetector` (P1) + `rally_gate` (already existed) ? no
reinvention.
68
69     **Status.** Design via the `fusion-p2-design` workflow ? implementation ? adversarial-
review workflow.
70     Lift to be **measured offline on the 6 golden videos via `experiment.compare(CONTROL,
fusion)`** ; the
71     held-shuttle-drop confirmation is GPU/owner-side. *(Numbers appended here when measured
? no claims before.)*
72
73     ### Person cue extracted to a reusable producer (fusion P1) ?
[#115](https://github.com/avidullu/badminton-highlight-indexer/pull/115), 2026-06-13
74     Lifted the torchvision detector + ByteTrack + velocity windowing out of `yolo_hybrid`
into
75     `backend/pipeline/detectors/person_detector.py::PersonDetector`. **No behavior change**
(proven by the
76     unchanged `process_video` output tests); `yolo_hybrid` now orchestrates only. **Why it's
a quality move:**
77     it makes "add the person cue to fusion" an `add-a-producer`, not a segmenter rewrite.
Suite 440?441 green.
78
79     ### Experiment harness: A/B without rollback risk ?
[#114](https://github.com/avidullu/badminton-highlight-indexer/pull/114), 2026-06-13
80     Shipped `backend/eval/experiment.py` ? the thin variant/experiment layer (the audit's
"-80% already
81     exists" held; pure glue over
`metrics`/`calibration`/`classifier`/`training`/`regression.gt_hash`). Gives
82     `Variant` (+ pinned `CONTROL`), `compare()` (labeled A/B, LOVO-honest),
**`compare_unlabeled()`** (the SxS
83     read on NEW unlabeled footage ? agreed / A-only / B-only), and a JSON experiment
leaderboard. **This is
84     the instrument every iteration is measured on.** 19 tests incl. the no-regression
invariant (all cues OFF
85     ? byte-identical to control). Suite 421?440 green.
86
87     ---
88
89     ## Archived predecessors (the historical trail)
90     Completed iterations live under [`archives/quality-iterations/`](archives/quality-
iterations/README.md):
91     - **`2026-06-initial/`** ? seed of the quality phase.
92     - **`2026-06-human-gt-tuning/`** ? first eval vs **human** GT on a partial-label video;
partial-label
93     harness + labeled-window methodology; baseline 0.650 ? tuned 0.742 (*fit-on-self*);
found+fixed a
94     negative-start ingestion bug (collector PR #13).
95     - **`2026-06-multivideo-lovo/`** ? the **n=6 corpus + honest LOVO** ; established the
heuristic is

```

```

96     **footage-fragile** (multi-court Kushagra 0.16 / gBaaddy 0.28), the ***"contrast"
metric** predicts
97     heuristic F1 monotonically, and a **learned per-frame prototype overtook the heuristic
at n=3 and holds
98     ~+0.10** (n=6: learned 0.583 vs heuristic 0.480), with **AUC 0.83?0.92 decoupled from
contrast**. This
99     is the data-backed green light for the owned model.
100
101     ---
102
103     ## Standing lessons (carry-forward)
104     - **Persist detection once, re-score variants forever.** Expensive/risky = per-frame
detection (GPU/WSL);
105     cheap/swappable = the decision. Most experiments never enter the served pipeline.
106     - **LOVO or it didn't happen.** Score on a held-out *video*, never a held-out window
from the same video.
107     Group by **match**, not file. Sign test needs n ? ~n=10/9-wins for p<0.05.
108     - **The multi-court confound is TEMPORAL, not spatial/loudness** ? three heuristics were
tested and failed
109     for the same reason (foreground spatial gate **?0.04**; audio energy **AUC ~0.6**;
court-region
110     separability **12?14% out-of-region**): the dead-time confusion is rally-vs-non-rally
activity in the
111     *same place and frequency band*. Only the learned model's temporal/structural features
crack it (AUC
112     0.83 on Kushagra/gBaaddy). **Don't re-derive these dead ends.**
113     - **Court-masking ? multi-court fix** (see above) ? it's a **precision** tool for off-
court persons and
114     one-sided feeds, a *different* lever; measure its lift, don't assume it.
115     - **Net-crossings are a strong rally signal** ? 97% of real rallies cross ?2x (median
5); single feeds ~1x.
116     Gate A leans on this.
117     - **Small-N feature discipline.** Few, scale/translation-invariant features
(counts/ratios, not raw
118     coords) + a logistic regression, not a net ? raw coords memorize this court/camera.
119     - **Calibration, not representation, is the small-N wall** ? learned probabilities
separate (high AUC) but
120     the operating point doesn't transfer with few training videos. More distinct matches
is the highest-ROI
121     unblock.
122     - **Gate A already existed, just wasn't wired to production** ? the cheapest precision
win is plumbing, not
123     new capability. Always check what the eval drivers already do that the live path
doesn't.
124

```

9. user (2026-06-13T19:17:33.223Z)

```

1     ""Multi-signal fusion segmenter (MULTI_SIGNAL_FUSION_PLAN P2).
2
3     Shuttle stays PRIMARY: this reuses an existing trajectory substrate (WASB by default,
4     TrackNet optional) to seed candidate windows exactly as `wasb_hybrid`/`tracknet_hybrid`
5     do ? it subclasses the same `TrajectoryHybridSegmenter` engine. It then ENRICHES each
6     shuttle-seeded window via two overridable seams the engine exposes:
7
8     - ``_enrich_candidate_stats`` ? always adds the **net-crossing count** (Gate A signal,
9     GPU-free, computed from the trajectory we already have); and, only when the person cue
10    is explicitly enabled (``indexing.fusion.cue_flags.person``), the person-cue features
11    (`fusion_features`), persisting them into ``candidate_segments.stats`` for the learned
12    fusion (P3). The person path lazily builds ONE `PersonDetector` ? so the default path
13    never imports torch / touches the GPU.
14    - ``_select_for_handoff`` ? optional served Gate A/B: when
``indexing.fusion.min_crossings``
15    (and, with the person cue, ``min_players``) are raised above 0, sub-threshold windows
are

```

```

16         dropped from the AI handoff (the held-shuttle / one-sided-feed false-positive fix).
17         **Persisted candidates remain the full superset** regardless, so the offline
experiment
18         harness can re-score any variant.
19
20     Defaults reproduce the substrate exactly: net_crossings/person features are persisted
but
21     nothing is gated (thresholds 0, person OFF), so `FusionSegmenter` with default config
seeds
22     and hands off identically to `wasb_hybrid`. Registered as ``fusion_hybrid``, NOT the
default
23     segmenter ? it is opt-in. Promotion to default happens only on measured LOVO lift
through the
24     experiment harness (EXPERIMENT_HARNESS.md), never silently.
25     """
26     from typing import Any, Dict, List, Optional, Tuple
27
28     from backend.pipeline.segmenters.base import segmenter_registry
29     from backend.pipeline.segmenters.trajectory_hybrid import TrajectoryHybridSegmenter
30     from backend.pipeline.detectors.base import DetectorRunner
31     from backend.pipeline.detectors.wasb_runner import WasbRunner, WasbConfig
32     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner, TrackNetConfig
33     from backend.pipeline.detectors import rally_gate
34     from backend.pipeline.detectors import fusion_features as ff
35
36
37     class FusionSegmenter(TrajectoryHybridSegmenter):
38         """Shuttle trajectory (primary) + net-crossing Gate A + optional person cue (Gate
B)."""
39
40         family = "fusion"
41         display_label = "Fusion"
42
43         def __init__(self, db: Any, config: Any):
44             super().__init__(db, config)
45             # Built lazily only if the person cue is enabled (keeps torch/GPU off the
default path).
46             self._person: Optional[Any] = None
47
48             @property
49             def name(self) -> str:
50                 return "fusion_hybrid"
51
52             @property
53             def description(self) -> str:
54                 return ("Multi-signal fusion: shuttle trajectory (WASB/TrackNet) seeds candidate
rallies, "
55                         "net-crossing Gate A + optional person-occupancy Gate B adjudicate them,
and the "
56                         "AI provider sets semantic boundaries.")
57
58             def _build_runner(self, idx_cfg: Dict[str, Any]) -> Tuple[DetectorRunner, Dict[str,
Any]]:
59                 substrate = idx_cfg.get("fusion", {}).get("substrate", "wasb")
60                 if substrate == "tracknet":
61                     cfg = TrackNetConfig.from_indexing_cfg(idx_cfg)
62                     return TrackNetRunner(cfg), {
63                         "weights_path": cfg.weights_path,
64                         "queue_length": cfg.queue_length,
65                         "fusion_substrate": "tracknet",
66                     }
67                 wcfg = WasbConfig.from_indexing_cfg(idx_cfg)
68                 return WasbRunner(wcfg), {"weights_path": wcfg.weights_path, "fusion_substrate":
"wasb"}
69

```

```

70     def _enrich_candidate_stats(self, cw: Dict[str, Any], points: List[Any], fps: float,
71                                frame_width: float, video_path: str,
72                                idx_cfg: Dict[str, Any]) -> Dict[str, Any]:
73         stats = super()._enrich_candidate_stats(cw, points, fps, frame_width,
video_path, idx_cfg)
74         # Gate-A signal ? net-crossing count for this window. Pure, GPU-free, always
persisted
75         # (single service feed crosses ~once; a real rally crosses >=2x). axis/net auto-
estimated
76         # over the whole trajectory by rally_gate (camera-agnostic).
77         gated = rally_gate.gate_windows(points, [[cw["start"], cw["end"]]], fps,
min_crossings=0)
78         stats["net_crossings"] = float(gated[0].crossings) if gated else 0.0
79
80         fcfg = idx_cfg.get("fusion", {})
81         if fcfg.get("cue_flags", {}).get("person", False):
82             stats.update(self._person_features(cw, points, fps, frame_width, video_path,
fcfg))
83         return stats
84
85     def _person_features(self, cw: Dict[str, Any], points: List[Any], fps: float,
86                         frame_width: float, video_path: str,
87                         fcfg: Dict[str, Any]) -> Dict[str, float]:
88         """GPU path (person cue ON): run the person detector over this window's
ORIGINAL-video
89         frame range and compute the scale/translation-invariant person features. Lazily
builds
90         ONE PersonDetector. ? real-video verification is owner-side; this wiring is
mock-tested."""
91         # Lazy import so the default (person-OFF) path never pulls torch/cv2 through
this module.
92         from backend.pipeline.detectors.person_detector import PersonDetector
93         if self._person is None:
94             self._person = PersonDetector(self.config)
95
96         frame_skip = self.config.get("indexing", {}).get("yolo_frame_skip", 3)
97         start_f, end_f = int(cw["start"] * fps), int(cw["end"] * fps)
98         det = self._person.detections_per_frame(video_path, start_f, end_f, frame_skip)
99         fw = det["frame_width"] or frame_width
100        fh = det["frame_height"] or 0.0
101
102        poly = fcfg.get("court_polygon")
103        polygon = [(float(p[0]), float(p[1])) for p in poly] if poly else None
104        net = (fcfg.get("net_axis", "y"), float(fcfg.get("net_position", 0.5)))
105        shuttle_in = [p for p in points if start_f <= p.frame <= end_f]
106        return ff.compute_person_features(det["frames"], shuttle_in, fw, fh,
107                                         court_polygon=polygon, net=net)
108
109     def _select_for_handoff(self, windows: List[Dict[str, Any]],
110                           window_stats: List[Dict[str, Any]],
111                           idx_cfg: Dict[str, Any]) -> List[int]:
112         """Served Gate A/B over the (already-enriched) per-window stats. Default
thresholds
113         (0) keep every window ? identical to the substrate. Persisted candidates are
114         untouched (full superset)."""
115         fcfg = idx_cfg.get("fusion", {})
116         min_crossings = int(fcfg.get("min_crossings", 0))
117         min_players = int(fcfg.get("min_players", 0))
118         person_on = bool(fcfg.get("cue_flags", {}).get("person", False))
119
120         keep: List[int] = []
121         for i, st in enumerate(window_stats):
122             if min_crossings > 0 and st.get("net_crossings", 0.0) < min_crossings:
123                 continue
124             if person_on and min_players > 0 and st.get("players_in_court", 0.0) <

```

```

min_players:
125             continue
126             keep.append(i)
127             return keep
128
129
130     segmenter_registry.register(FusionSegmenter)
131

```

10. assistant (2026-06-13T19:17:37.073Z)

Now the segmenter engine, the feature module, the person detector, the database substrate, and the docs index.

11. user (2026-06-13T19:17:37.563Z)

```

1      """Shared engine for trajectory-substrate hybrid segmenters.
2
3      tracknet_hybrid and wasb_hybrid were copy-paste twins ? CODE_MAP carried a
4      standing footgun: "any P3/P4 fix must be applied to BOTH". This base class is
5      now the single home for the whole pipeline shape:
6
7          healthcheck gate -> P2 trajectory -> candidate windows (+ persistence)
8          -> P3 AI-provider handoff over padded clips -> P4 DB population/linking
9
10     A concrete subclass supplies only its identity and runner wiring:
11     - ``family`` ? config section under ``indexing.*``, output subdir, and
12       the ``<family>-hybrid-<model>`` detector tag in metadata.
13     - ``display_label`` ? human label for log lines ("WASB", "TrackNet").
14     - ``name`` / ``description`` properties (the registry contract).
15     - ``_build_runner(idx_cfg)`` ? returns (DetectorRunner, detector-specific
16       detection_params extras such as weights_path).
17
18     This base is deliberately NOT registered; only concrete subclasses register.
19     """
20     import os
21     import time
22     import logging
23     import concurrent.futures
24     from typing import Any, Dict, List, Optional, Tuple
25
26     import cv2
27
28     from backend.pipeline.segmenters.base import BasePlaySegmenter
29     from backend.utils.video import get_video_duration, extract_video_chunk
30     from backend.sports import sport_registry
31     from backend.providers import provider_registry
32     from backend.pipeline.detectors.base import DetectorRunner
33
34     logger = logging.getLogger(__name__)
35
36
37     def _fmt_ts(seconds: float) -> str:
38         seconds = max(0, int(seconds))
39         h, rem = divmod(seconds, 3600)
40         m, s = divmod(rem, 60)
41         return f"{h:02d}:{m:02d}:{s:02d}"
42
43
44     class TrajectoryHybridSegmenter(BasePlaySegmenter):
45         """Trajectory-detector hybrid: precise candidate rallies + AI semantic
46         boundaries."""
47
48         #: config section under `indexing` (velocity_threshold lives there), the
49         #: output subdir for trajectories, and the metadata detector-tag prefix.

```

```

49     family: str = ""
50     #: human-readable label used in log lines.
51     display_label: str = ""
52
53     def _build_runner(self, idx_cfg: Dict[str, Any]) -> Tuple[DetectorRunner, Dict[str,
Any]]:
54         """Return (runner, detector-specific detection_params extras)."""
55         raise NotImplementedError
56
57     @staticmethod
58     def _probe_fps_width(video_path: str) -> tuple:
59         """Return (fps, frame_width) for a video, with sensible fallbacks."""
60         cap = cv2.VideoCapture(video_path)
61         fps = cap.get(cv2.CAP_PROP_FPS) if cap.isOpened() else 0.0
62         width = cap.get(cv2.CAP_PROP_FRAME_WIDTH) if cap.isOpened() else 0.0
63         cap.release()
64         return (fps if fps > 0 else 30.0, width if width > 0 else 1280.0)
65
66     @staticmethod
67     def _shot_count_from(segment: Dict[str, Any], duration: float) -> int:
68         """Provider-reported exchange count, else a duration-based estimate.
69
70         One canonical implementation ? the twins had drifted (wasb relied on
71         ``int(None)`` raising into the fallback; same result, by accident).
72         """
73         shot_count = segment.get("shuttle_exchange_count")
74         if shot_count is None:
75             return max(3, int(duration / 0.8))
76         try:
77             return int(shot_count)
78         except (ValueError, TypeError):
79             return max(3, int(duration / 0.8))
80
81     # --- Fusion seams (identity by default ? wasb_hybrid/tracknet_hybrid unchanged)
82     -----
83     def _enrich_candidate_stats(self, cw: Dict[str, Any], points: List[Any], fps: float,
84                                frame_width: float, video_path: str,
85                                idx_cfg: Dict[str, Any]) -> Dict[str, Any]:
86         """Stats dict persisted into ``candidate_segments.stats`` for one window. The
87         BASE
88         returns exactly the 4 motion keys, so the trajectory twins persist byte-
89         identical
90         rows; ``FusionSegmenter`` overrides this to add ``net_crossings`` (+ person-cue
91         features). ``points`` are the whole-video trajectory points
92         (TrajectoryPoint)."""
93         return {
94             "peak_velocity": cw["peak_velocity"],
95             "mean_velocity": cw["mean_velocity"],
96             "active_frame_count": cw["active_frame_count"],
97             "track_id_count": cw["track_id_count"],
98         }
99
100     def _select_for_handoff(self, windows: List[Dict[str, Any]],
101                            window_stats: List[Dict[str, Any]],
102                            idx_cfg: Dict[str, Any]) -> List[int]:
103         """Indices of the candidate windows that proceed to the AI handoff (and thus may
104         become play_segments). The BASE sends ALL windows (no gating ? twins unchanged);
105         ``FusionSegmenter`` overrides this to apply Gate A/B when thresholds are raised.
106         Persisted candidates are NOT affected ? they remain the full superset for
107         offline
108         re-scoring."""
109         return list(range(len(windows)))
110
111     def process_video(self, video_id: str, video_path: str, # type: ignore[override]
112                      player_pool: Optional[List[str]] = None,

```

```

108         max_frames: Optional[int] = None) -> List[Dict[str, Any]]:
109     # override-ignore: the ABC declares the Result contract (Success|Failure)
110     # but every live segmenter still returns a bare list ? the documented
111     # deliberate-incremental gap (CODE_MAP gotchas; main.py handles both).
112     label = self.display_label
113     # --- Sport profile + AI provider wiring (identical across segmenters) ---
114     sport_name = self.config.get("sport", "badminton")
115     try:
116         sport_profile = sport_registry.get(sport_name)
117     except KeyError as e:
118         logger.error(str(e))
119         return []
120
121     idx_cfg = self.config.get("indexing", {})
122     provider_name = self.config.get("ai_provider", idx_cfg.get("ai_provider",
123 "gemini"))
124     model_name = self.config.get("ai_model", idx_cfg.get("ai_model",
125 "gemini-2.5-flash"))
126
127     api_key_env_var = f"{provider_name.upper()}_API_KEY"
128     api_key = os.environ.get(api_key_env_var)
129     if not api_key:
130         logger.error(
131             f"{api_key_env_var} environment variable is not set! "
132             f"To run the {provider_name} handoff, set your API key. "
133             f"In PowerShell: $env:{api_key_env_var}=\"your_api_key_here\"")
134         return []
135     try:
136         provider = provider_registry.get(provider_name, api_key=api_key,
137 model_name=model_name)
138     except KeyError as e:
139         logger.error(str(e))
140         return []
141
142     video_duration = get_video_duration(video_path)
143     if video_duration <= 0:
144         logger.error("Invalid video duration.")
145         return []
146     fps, frame_width = self._probe_fps_width(video_path)
147
148     # --- Runner config + windowing thresholds ---
149     runner, runner_params = self._build_runner(idx_cfg)
150
151     velocity_thresh = idx_cfg.get(self.family, {}).get("velocity_threshold", 0.02)
152     merge_gap = idx_cfg.get("dead_time_merge_gap", 2.0)
153     min_window_duration = idx_cfg.get("min_action_window_duration", 2.0)
154     handoff_padding = idx_cfg.get("ai_handoff_padding", 2.0)
155     ai_max_workers = idx_cfg.get("ai_max_workers", 5)
156     log_candidates = idx_cfg.get("log_yolo_candidates", True)
157     store_candidates = idx_cfg.get("store_yolo_candidates", True)
158
159     detection_params = {
160         "detector": self.name,
161         "velocity_thresh": velocity_thresh,
162         "merge_gap": merge_gap,
163         "min_action_window_duration": min_window_duration,
164         **runner_params,
165     }
166
167     # --- Phase 1: env healthcheck (fail fast with a clear message) ---
168     ok, msg = runner.healthcheck()
169     if not ok:
170         logger.error("%s WSL environment not ready: %s", label, msg)
171         return []

```

```

170         logger.info("%s WSL env OK: %s", label, msg)
171
172         # --- Phase 2: trajectory -> candidate rally windows ---
173         # Trajectory detectors process the whole video in one pass (they carry
174         # their own frame buffering), so unlike yolo_hybrid we don't pre-chunk.
175         t_start = time.time()
176         out_dir = os.path.join("output", self.family)
177         csv_path = runner.run_predict(video_path, out_dir, video_id=video_id)
178         if not csv_path:
179             logger.error("%s produced no trajectory; aborting.", label)
180             return []
181
182         points = runner.parse_trajectory_csv(csv_path)
183         logger.info("Parsed %d trajectory points (%d visible).",
184                     len(points), sum(1 for p in points if p.visible))
185
186         all_candidate_windows = runner.trajectory_to_action_windows(
187             points, fps=fps, frame_width=frame_width, chunk_start=0.0,
188             velocity_thresh=velocity_thresh, merge_gap=merge_gap,
189             min_window_duration=min_window_duration, chunk_index=0,
190         )
191         logger.info("Phase 2 (%s) completed in %.1fs. Candidate windows: %d",
192                     label, time.time() - t_start, len(all_candidate_windows))
193
194         if log_candidates:
195             for cw in all_candidate_windows:
196                 logger.info(f"[{label}] rally
197 {fmt_ts(cw['start'])}-{fmt_ts(cw['end'])} "
198                             f"({cw['end'] - cw['start']:.1f}s) |
199 frames={cw['active_frame_count']} "
200                             f"peak_v={cw['peak_velocity']:.3f}
201 mean_v={cw['mean_velocity']:.3f}")
202
203         # Per-window stats via the enrichment hook ? base = the 4 motion keys; the
204         fusion
205         # segmenter adds net_crossings + person-cue features. Computed once, reused for
206         both
207         # persistence and the handoff gate below.
208         window_stats = [
209             self._enrich_candidate_stats(cw, points, fps, frame_width, video_path,
210             idx_cfg)
211             for cw in all_candidate_windows
212         ]
213
214         # Persist raw candidates for the fine-tuning dataset (same table as YOLO).
215         candidate_ids: List[Optional[int]] = [None] * len(all_candidate_windows)
216         if store_candidates:
217             for i, cw in enumerate(all_candidate_windows):
218                 candidate_ids[i] = self.db.add_candidate_segment(
219                     video_id, detector=self.name,
220                     start_time=cw["start"], end_time=cw["end"],
221                     chunk_index=cw["chunk_index"], confidence=min(1.0,
222                     cw["peak_velocity"]),
223                     stats=window_stats[i], detection_params=detection_params,
224                 )
225
226         if not all_candidate_windows:
227             return []
228
229         # --- Phase 3: AI provider handoff over padded candidate clips ---
230         # Which windows reach the handoff (base = all; fusion may apply Gate A/B).
231         Persisted
232         # candidates above stay the full superset ? only the served play_segments are
233         gated.
234         handoff_indices = self._select_for_handoff(all_candidate_windows, window_stats,

```



```

idx_cfg)
226         handoff_dir = os.path.join("output", f"chunks_{provider_name}_handoff")
227         os.makedirs(handoff_dir, exist_ok=True)
228         logger.info("Phase 3: %s validation on %d candidate windows...",
229                     provider_name, len(handoff_indices))
230
231     def process_candidate_window(wi: int, window: Dict[str, Any]) -> List[dict]:
232         w_start = max(0, window["start"] - handoff_padding)
233         w_end = min(video_duration, window["end"] + handoff_padding)
234         w_dur = w_end - w_start
235         w_path = os.path.join(handoff_dir, f"handoff_{video_id}_{wi}.mp4")
236         if not extract_video_chunk(video_path, w_start, w_dur, w_path,
reencode=True):
237             return []
238         prompt = sport_profile.get_prompt(chunk_duration=w_dur)
239         segments, _ = provider.analyze_video(w_path, prompt, chunk_duration=w_dur,
240                                             label=f"Handoff {wi+1}")
241
242         valid = []
243         for seg in segments:
244             try:
245                 seg["start_time"] = float(seg["start_time"]) + w_start
246                 seg["end_time"] = float(seg["end_time"]) + w_start
247                 seg["_candidate_id"] = candidate_ids[wi]
248                 valid.append(seg)
249             except (ValueError, TypeError, KeyError) as e:
250                 logger.warning("Dropping segment with unparseable timestamps: %s ?
%s", seg, e)
251
252         try:
253             os.remove(w_path)
254         except OSError:
255             pass
256         return valid
257
258     ai_segments: List[dict] = []
259     with concurrent.futures.ThreadPoolExecutor(max_workers=ai_max_workers) as
executor:
260         futures = [executor.submit(process_candidate_window, wi,
all_candidate_windows[wi])
261                     for wi in handoff_indices]
262         for future in futures:
263             ai_segments.extend(future.result())
264
265     try:
266         os.rmdir(handoff_dir)
267     except OSError:
268         pass
269     logger.info("Phase 3 completed. AI confirmed %d semantic play segments.",
len(ai_segments))
270
271     # --- Phase 4: DB population + candidate linking ---
272     segments_data = []
273     for segment in ai_segments:
274         try:
275             start = float(segment["start_time"])
276             end = float(segment["end_time"])
277             dur = end - start
278             metadata = {
279                 "shot_count": self._shot_count_from(segment, dur),
280                 "shot_count_confidence": segment.get("shot_count_confidence",
"Unknown"),
281                 "detector": f"{self.family}-hybrid-{model_name}",
282             }
283             segment_id = self.db.add_play_segment(video_id, start, end, None, None,
metadata)
284
285             if segment_id:
286                 candidate_id = segment.get("_candidate_id")

```

```

283         if candidate_id is not None:
284             self.db.link_candidate_to_segment(candidate_id, segment_id)
285             segments_data.append({
286                 "id": segment_id, "start_time": start, "end_time": end,
287                 "duration": dur, "metadata": metadata,
288             })
289         except Exception as e:
290             logger.warning("Skipping invalid segment struct: %s", e)
291
292     return segments_data
293

```

12. user (2026-06-13T19:17:38.028Z)

```

1     """Pure per-window person-cue features for multi-signal rally fusion
(MULTI_SIGNAL_FUSION_PLAN §3.2/§4).
2
3     NO I/O, no torch, no cv2 ? mirrors `rally_gate.py`'s discipline so the whole feature
4     math is unit-testable in a GPU-free dev container. The detection (person boxes,
5     shuttle trajectory) is produced elsewhere (`PersonDetector`, the WASB runner); this
6     module only turns those signals into a handful of scale/translation-invariant
7     features (counts / ratios / zone-membership ? never raw pixel coordinates, which would
8     memorise this court/camera and fail to generalise off a small labelled set).
9
10    Conventions:
11    - A person box is `(x1, y1, x2, y2)` in pixels; `(track_id, box)` when motion
needs it.
12    - `per_frame` is `{frame_idx: [(track_id, box), ...]}` keyed by the ORIGINAL video
frame
13        index (so it aligns directly with the shuttle trajectory's `frame` ? no clip re-
timing).
14    - A shuttle point has `frame`, `visible`, `x`, `y` (pixels) ? the WASB
trajectory shape.
15    - `court_polygon` is normalised 0-1 `(x, y), ...)` or `None` (None ? no mask,
every
16        detection counts ? the player-count-only mode).
17    - `net` is `(axis 'x'|'y', position 0-1 normalised)` or `None` (None ? two-sided =
0.0).
18
19    Every function degrades gracefully (returns 0.0 / empty) on missing inputs; outputs are
20    floats in [0, 1] except `mean_player_motion` (a non-negative normalised speed).
21    """
22    from __future__ import annotations
23
24    from typing import Dict, List, Optional, Sequence, Tuple
25
26    BBox = Tuple[float, float, float, float]
27    Net = Tuple[str, float]
28
29
30    def point_in_polygon(point: Tuple[float, float], polygon: Sequence[Tuple[float, float]])
-> bool:
31        """Ray-casting point-in-polygon. `point` and `polygon` share a coordinate space
32        (we use normalised 0-1). Empty/degenerate polygon (<3 vertices) ? False."""
33        if polygon is None or len(polygon) < 3:
34            return False
35        x, y = point
36        inside = False
37        n = len(polygon)
38        j = n - 1
39        for i in range(n):
40            xi, yi = polygon[i]
41            xj, yj = polygon[j]
42            # Does the horizontal ray at y cross edge (i, j)?
43            if (yi > y) != (yj > y):

```

```

44         x_cross = (xj - xi) * (y - yi) / (yj - yi) + xi if yj != yi else xi
45         if x < x_cross:
46             inside = not inside
47         j = i
48     return inside
49
50
51     def _foot_norm(box: BBox, frame_width: float, frame_height: float) -> Tuple[float,
float]:
52         """Normalised foot point (bottom-centre) of a person box ? where the player stands,
53         the right anchor for court membership. Returns (0,0) if frame dims are unusable."""
54         if frame_width <= 0 or frame_height <= 0:
55             return (0.0, 0.0)
56         x1, _y1, x2, y2 = box
57         return ((x1 + x2) / 2.0 / frame_width, y2 / frame_height)
58
59
60     def _center_norm(box: BBox, frame_width: float, frame_height: float) -> Tuple[float,
float]:
61         if frame_width <= 0 or frame_height <= 0:
62             return (0.0, 0.0)
63         x1, y1, x2, y2 = box
64         return ((x1 + x2) / 2.0 / frame_width, (y1 + y2) / 2.0 / frame_height)
65
66
67     def person_in_court(box: BBox, frame_width: float, frame_height: float,
68                        court_polygon: Optional[Sequence[Tuple[float, float]]]) -> bool:
69         """Is this person on the court? Foot-point-in-polygon when a polygon is supplied;
70         True (count everyone) when it is None."""
71         if court_polygon is None:
72             return True
73         return point_in_polygon(_foot_norm(box, frame_width, frame_height), court_polygon)
74
75
76     def in_court_counts(per_frame: Dict[int, List[Tuple[int, BBox]]], frame_width: float,
77                        frame_height: float,
78                        court_polygon: Optional[Sequence[Tuple[float, float]]]) ->
List[int]:
79         """Per-frame count of in-court persons (one entry per sampled frame)."""
80         return [
81             sum(1 for _tid, box in dets if person_in_court(box, frame_width, frame_height,
court_polygon))
82             for _frame_idx, dets in sorted(per_frame.items())
83         ]
84
85
86     def _median(vals: Sequence[float]) -> float:
87         if not vals:
88             return 0.0
89         s = sorted(vals)
90         n = len(s)
91         return float(s[n // 2] if n % 2 else 0.5 * (s[n // 2 - 1] + s[n // 2]))
92
93
94     def players_in_court(counts: Sequence[int]) -> float:
95         """Median per-frame in-court count (?2 singles / 4 doubles; 0-1 = dead time)."""
96         return _median(list(counts))
97
98
99     def in_court_ratio(counts: Sequence[int], min_players: int = 2) -> float:
100         """Fraction of sampled frames with >= `min_players` in court (track stability vs
flicker)."""
101         if not counts:
102             return 0.0
103         return sum(1 for c in counts if c >= min_players) / len(counts)

```

```

104
105
106 def mean_player_motion(per_frame: Dict[int, List[Tuple[int, BBox]]],
107                        frame_width: float, frame_height: float) -> float:
108     """Mean per-track centre displacement between consecutive sampled frames, normalised
109     by the frame width (resolution-invariant). 0.0 if <2 frames or no shared tracks."""
110     if frame_width <= 0 or len(per_frame) < 2:
111         return 0.0
112     frames = sorted(per_frame.keys())
113     steps: List[float] = []
114     for a, b in zip(frames, frames[1:]):
115         prev = {tid: box for tid, box in per_frame[a]}
116         for tid, box in per_frame[b]:
117             if tid in prev:
118                 cx0, cy0 = _center_norm(prev[tid], frame_width, frame_height)
119                 cx1, cy1 = _center_norm(box, frame_width, frame_height)
120                 steps.append(((cx1 - cx0) ** 2 + (cy1 - cy0) ** 2) ** 0.5)
121     if not steps:
122         return 0.0
123     return sum(steps) / len(steps)
124
125
126 def two_sided_motion(per_frame: Dict[int, List[Tuple[int, BBox]]], frame_width: float,
127                    frame_height: float, net: Optional[Net],
128                    court_polygon: Optional[Sequence[Tuple[float, float]]]) -> float:
129     """Fraction of sampled frames with >=1 in-court person on BOTH net halves.
130
131     A real rally is two-sided; a single player feeding/holding the shuttle is one-sided.
132     The net splits the play axis at `net=(axis, position)` in normalised coords. None ?
133     0.0.
134     """
135     if net is None or not per_frame:
136         return 0.0
137     axis, pos = net
138     both = 0
139     for _frame_idx, dets in per_frame.items():
140         near, far = False, False
141         for _tid, box in dets:
142             if not person_in_court(box, frame_width, frame_height, court_polygon):
143                 continue
144             fx, fy = _foot_norm(box, frame_width, frame_height)
145             coord = fx if axis == "x" else fy
146             if coord < pos:
147                 near = True
148             else:
149                 far = True
150             if near and far:
151                 both += 1
152     return both / len(per_frame)
153
154 def _shuttle_in_box(sx: float, sy: float, box: BBox, frame_width: float, frame_height:
float,
155                    pad_frac: float) -> bool:
156     """Is the (pixel) shuttle point inside a person box, expanded by `pad_frac` of frame
157     width/height (the 'adjacent / in-hand' tolerance)?"
158     x1, y1, x2, y2 = box
159     px = pad_frac * frame_width
160     py = pad_frac * frame_height
161     return (x1 - px) <= sx <= (x2 + px) and (y1 - py) <= sy <= (y2 + py)
162
163
164 def shuttle_player_colocation(per_frame: Dict[int, List[Tuple[int, BBox]]],
shuttle_points,
165                               frame_width: float, frame_height: float,

```

```

166         pad_frac: float = 0.03) -> float:
167     """Fraction of (visible shuttle frames that were also person-sampled) in which the
168     shuttle sits inside/adjacent to a person box. **held/in-hand ? high; in-flight ?
169     ~0**
170     ? the direct held-shuttle discriminator. 0.0 if no aligned visible shuttle frames.
171     """
172     if not per_frame or frame_width <= 0 or frame_height <= 0:
173         return 0.0
174     aligned = 0
175     colocated = 0
176     for p in shuttle_points or []:
177         if not getattr(p, "visible", False):
178             continue
179         dets = per_frame.get(int(p.frame))
180         if dets is None: # shuttle frame wasn't person-sampled ? can't judge
181             continue
182         aligned += 1
183         if any(_shuttle_in_box(p.x, p.y, box, frame_width, frame_height, pad_frac)
184             for _tid, box in dets):
185             colocated += 1
186     return colocated / aligned if aligned else 0.0
187
188 # The ~5 person features persisted into candidate `stats` (order = the learned-fusion
189 # feature columns appended after the shuttle features in distill_local.FEATURE_NAMES).
190 PERSON_FEATURE_NAMES = [
191     "players_in_court", "two_sided_motion", "mean_player_motion",
192     "in_court_ratio", "shuttle_player_colocation",
193 ]
194
195
196 def compute_person_features(per_frame: Dict[int, List[Tuple[int, BBox]]],
197 shuttle_points,
198                             frame_width: float, frame_height: float,
199                             court_polygon: Optional[Sequence[Tuple[float, float]]] =
200 None,
201                             net: Optional[Net] = None,
202                             min_players: int = 2) -> Dict[str, float]:
203     """Aggregate the per-window person features from per-frame detections + the in-
204     window
205     shuttle trajectory. Pure; returns exactly `PERSON_FEATURE_NAMES`, all floats."""
206     counts = in_court_counts(per_frame, frame_width, frame_height, court_polygon)
207     return {
208         "players_in_court": players_in_court(counts),
209         "two_sided_motion": two_sided_motion(per_frame, frame_width, frame_height, net,
210 court_polygon),
211         "mean_player_motion": mean_player_motion(per_frame, frame_width, frame_height),
212         "in_court_ratio": in_court_ratio(counts, min_players),
213         "shuttle_player_colocation": shuttle_player_colocation(
214             per_frame, shuttle_points, frame_width, frame_height),
215     }
216

```

13. assistant (2026-06-13T19:17:41.803Z)

Now the person detector's `detections\_per\_frame`, the database substrate (events + candidate segments), the docs index, and rally_gate for the net estimate.

14. user (2026-06-13T19:17:42.757Z)

```

1     """Person-detection rally cue ? torchvision detector + ByteTrack + motion windowing.
2
3     Extracted verbatim from `yolo_hybrid` (fusion P1, `docs/MULTI_SIGNAL_FUSION_PLAN.md`):
4     the **person cue is a swappable producer**, not logic baked into one segmenter. Today
5     `HybridYoloSegmenter` is the only consumer (it delegates Stage-1 here, byte-

```

```

identically);
6     next the fusion segmenter consumes the SAME producer over shuttle-seeded windows, and
7     later the learned classifier reads its persisted features. Keeping detection here (not
in
8     the segmenter) is what makes "add a cue = register a producer" true.
9
10    Behaviour is unchanged from the in-lined version ? only the home moved:
11    - `lazy_load_model`           ? load a permissively-licensed torchvision detector
12    - `make_tracker`             ? a fresh supervision ByteTrack per chunk
13    - `detect_and_track`         ? one frame ? confident persons with persistent IDs
14    - `extract_action_windows_from_chunk` ? a chunk ? high-motion candidate windows
15        with the motion stats (`peak/mean_velocity`, `active_frame_count`, `track_id_count`)
16        that feed the candidate `stats` JSON and the learned fusion features.
17
18    Heavy deps (torchvision, supervision) are imported lazily inside the methods so
importing
19    the segmenter registry never pulls torch on a machine that lacks it (and tests can mock
20    the seams). Licensing: torchvision detectors are BSD + COCO weights; ByteTrack is MIT
21    (ADR-005; ultralytics' AGPL YOLO was dropped).
22    """
23    import logging
24    from typing import Any, Dict, List, Optional, Tuple
25
26    import cv2
27    import torch
28
29    from backend.utils.geometry import get_velocity_normalised
30
31    logger = logging.getLogger(__name__)
32
33    # torchvision detection models are trained on COCO where the "person" category is
34    # class 1 (index 0 is the implicit background). NB: ultralytics YOLO used class 0 ?
35    # this off-by-one is the classic gotcha when swapping detectors. Keep it named.
36    _PERSON_CLASS_ID = 1
37
38    # Permissively-licensed torchvision detectors selectable via config
39    # `indexing.detector_model`. All are BSD-licensed (code + COCO weights).
40    _DETECTOR_FACTORIES = {
41        "fasterrcnn_resnet50_fpn": "fasterrcnn_resnet50_fpn",
42        "fasterrcnn_mobilenet_v3_large_fpn": "fasterrcnn_mobilenet_v3_large_fpn",
43        "retinanet_resnet50_fpn": "retinanet_resnet50_fpn",
44    }
45    _DEFAULT_DETECTOR = "fasterrcnn_resnet50_fpn"
46
47
48    class PersonDetector:
49        """Stage-1 person detection + tracking + motion-window extraction (a rally cue).
50
51        Owns the torchvision model + per-chunk ByteTrack tracker. Constructed with the run
52        ``config`` (reads ``indexing.{use_gpu, detector_model, detector_score_threshold,
53        min_action_window_duration}``); the model loads lazily on first use.
54        """
55
56        def __init__(self, config: Dict[str, Any]):
57            self.config = config
58            # Dynamic torchvision module (None until lazy_load_model); typed Any so the
59            # `self.model([tensor])` inference call type-checks without a quarantine.
60            self.model: Any = None
61            self.device: Optional[str] = None
62
63        def lazy_load_model(self) -> None:
64            if self.model is not None:
65                return
66
67            # Imported lazily so the module imports without torchvision present and so

```

```

68         # the test suite (which pre-sets self.model) never pulls in model weights.
69         from torchvision.models import detection as tv_detection
70
71         idx_cfg = self.config.get("indexing", {})
72         use_gpu = idx_cfg.get("use_gpu", True)
73         self.device = "cuda" if use_gpu and torch.cuda.is_available() else "cpu"
74         if use_gpu and not torch.cuda.is_available():
75             logger.warning("use_gpu is True but CUDA is not available. Falling back to
CPU.")
76
77         model_name = idx_cfg.get("detector_model", _DEFAULT_DETECTOR)
78         if model_name not in _DETECTOR_FACTORIES:
79             logger.warning(f"Unknown detector_model '{model_name}'. Falling back to
{_DEFAULT_DETECTOR}.")
80             model_name = _DEFAULT_DETECTOR
81
82         builder = getattr(tv_detection, _DETECTOR_FACTORIES[model_name])
83         logger.info(f"Loading torchvision detector '{model_name}' on {self.device}...")
84         # weights="DEFAULT" pulls the best available COCO-pretrained checkpoint.
85         self.model = builder(weights="DEFAULT")
86         self.model.eval()
87         self.model.to(self.device)
88
89     def make_tracker(self, fps: float) -> Any:
90         """Create a fresh ByteTrack tracker for one chunk (track IDs are chunk-local).
91
92         Imported lazily so tests can mock this method without supervision installed.
93         """
94         import supervision as sv
95         frame_rate = max(1, int(round(fps)))
96         return sv.ByteTrack(frame_rate=frame_rate)
97
98     def detect_and_track(self, frame: Any, tracker: Any,
99                         score_thresh: float) -> List[Tuple[int, Tuple[float, float,
float, float]]]:
100         """Detect persons in a frame and assign persistent track IDs.
101
102         Returns a list of (track_id, bbox) where bbox is (x1, y1, x2, y2) in pixels.
103         This replaces the old ultralytics ``model.track(...)`` call, which bundled
104         detection AND association: torchvision does the detection, supervision's
105         ByteTrack does the association.
106         """
107         import supervision as sv
108
109         # cv2 frames are BGR uint8 HxWxC; torchvision wants RGB float [0,1] CxHxW.
110         rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
111         tensor = torch.from_numpy(rgb).permute(2, 0,
1).float().div(255.0).to(self.device)
112         with torch.no_grad():
113             outputs = self.model([tensor])
114
115         out = outputs[0]
116         boxes = out["boxes"].detach().cpu().numpy()
117         labels = out["labels"].detach().cpu().numpy()
118         scores = out["scores"].detach().cpu().numpy()
119
120         # Keep only confident person detections.
121         mask = (labels == _PERSON_CLASS_ID) & (scores >= score_thresh)
122         if not mask.any():
123             return []
124
125         # ByteTrack requires confidence to be populated (it filters on it internally).
126         dets = sv.Detections(xyxy=boxes[mask], confidence=scores[mask],
class_id=labels[mask])
127         tracked = tracker.update_with_detections(dets)

```

```

128
129     result: List[Tuple[int, Tuple[float, float, float, float]]] = []
130     for i in range(len(tracked)):
131         tid = tracked.tracker_id[i]
132         if tid is None:
133             continue
134         x1, y1, x2, y2 = tracked.xyxy[i]
135         result.append((int(tid), (float(x1), float(y1), float(x2), float(y2))))
136     return result
137
138     def extract_action_windows_from_chunk(self, chunk_path: str, chunk_start: float,
139                                         velocity_thresh: float, merge_gap: float,
140                                         frame_skip: int = 3,
141                                         chunk_index: Optional[int] = None) ->
List[Dict[str, Any]]:
142         """
143         Runs person detection + tracking on a video chunk and extracts high-motion
144         action windows. Returns a list of dicts in the full video's timeframe, each
with:
145             start, end (seconds), active_frame_count, peak_velocity, mean_velocity,
146             track_id_count, chunk_index.
147
148         Args:
149             frame_skip: Process every Nth frame (1 = every frame). Higher values
150                 dramatically reduce inference cost with minimal quality loss.
151             chunk_index: Index of the source chunk (recorded on each window for
traceability).
152         """
153         cap = cv2.VideoCapture(chunk_path)
154         if not cap.isOpened():
155             logger.error(f"Could not open {chunk_path}")
156             return []
157
158         fps = cap.get(cv2.CAP_PROP_FPS)
159         if fps <= 0:
160             fps = 30.0
161
162         frame_width = cap.get(cv2.CAP_PROP_FRAME_WIDTH)
163         if frame_width <= 0:
164             frame_width = 1280.0 # Sensible fallback
165
166         # Effective sampling rate after frame-skipping
167         effective_fps = fps / max(frame_skip, 1)
168
169         # A fresh tracker per chunk ? track IDs only need to be consistent within a
170         # chunk (windows are computed per chunk, then merged across chunks by time).
171         tracker = self.make_tracker(effective_fps)
172         score_thresh = self.config.get("indexing", {}).get("detector_score_threshold",
0.5)
173
174         # Per active frame we keep (frame_idx, peak_velocity, set_of_active_track_ids)
175         # so windows can be enriched with motion stats for logging + fine-tuning.
176         active_frames = []
177         frame_idx = 0
178         track_history: Dict[int, Tuple[float, float, float, float]] = {}
179
180         while True:
181             ret, frame = cap.read()
182             if not ret:
183                 break
184
185             # Skip frames for performance ? at 30fps, every 3rd frame still
186             # gives ~10 samples/second, plenty for human-scale velocity.
187             if frame_idx % max(frame_skip, 1) != 0:
188                 frame_idx += 1

```



```

189         continue
190
191     # Detect persons + assign persistent track IDs.
192     detections = self.detect_and_track(frame, tracker, score_thresh)
193
194     frame_peak_velocity = 0.0
195     frame_active_ids = set()
196     current_ids = set()
197
198     for track_id, bbox in detections:
199         current_ids.add(track_id)
200
201         if track_id in track_history:
202             prev_bbox = track_history[track_id]
203             velocity = get_velocity_normalised(prev_bbox, bbox, effective_fps,
frame_width)
204             if velocity > velocity_thresh:
205                 frame_peak_velocity = max(frame_peak_velocity, velocity)
206                 frame_active_ids.add(track_id)
207
208             track_history[track_id] = bbox
209
210     # Prune stale IDs that are no longer visible ? prevents false
211     # velocity spikes when the tracker reassigns an old ID.
212     track_history = {k: v for k, v in track_history.items() if k in current_ids}
213
214     if frame_active_ids:
215         active_frames.append((frame_idx, frame_peak_velocity, frame_active_ids))
216
217     frame_idx += 1
218
219     cap.release()
220
221     # Group active frames into continuous windows
222     if not active_frames:
223         return []
224
225     max_gap_frames = int(merge_gap * fps)
226
227     def _finalize(group: List[Tuple[int, float, set]]) -> Dict[str, Any]:
228         velocities = [v for _, v, _ in group]
229         track_ids = set().union(*[ids for _, _, ids in group])
230         return {
231             "start": group[0][0] / fps + chunk_start,
232             "end": group[-1][0] / fps + chunk_start,
233             "active_frame_count": len(group),
234             "peak_velocity": max(velocities),
235             "mean_velocity": sum(velocities) / len(velocities),
236             "track_id_count": len(track_ids),
237             "chunk_index": chunk_index,
238         }
239
240     windows = []
241     group = [active_frames[0]]
242     for af in active_frames[1:]:
243         if af[0] - group[-1][0] <= max_gap_frames:
244             group.append(af)
245         else:
246             windows.append(_finalize(group))
247             group = [af]
248     windows.append(_finalize(group))
249
250     # Filter out extremely short action windows
251     min_window_duration = self.config.get("indexing",
{}).get("min_action_window_duration", 2.0)

```

```

252         windows = [w for w in windows if (w["end"] - w["start"]) >= min_window_duration]
253
254     return windows
255
256     def detections_per_frame(self, video_path: str, start_frame: int, end_frame: int,
257                             frame_skip: int = 3) -> Dict[str, Any]:
258         """Per-frame person detections over `[start_frame, end_frame]` of the ORIGINAL
video.
259
260         Used by the fusion segmenter (MULTI_SIGNAL_FUSION_PLAN P2). Unlike
261         ``extract_action_windows_from_chunk`` (which detects ITS OWN windows and
discards the
262         per-frame boxes), this returns the raw per-frame detections so the fusion
features can
263         align them with the shuttle trajectory. It reads the source by **absolute frame
index**
264         (seek + sequential read), NOT a re-extracted clip, so the returned frame indices
match
265         the WASB trajectory's ``.frame`` exactly ? that alignment is what
266         ``fusion_features.shuttle_player_colocation`` / ``two_sided_motion`` rely on.
267
268         Returns ``{"frames": {frame_idx: [(track_id, bbox), ...]}, "fps", "frame_width",
269         "frame_height"}``. A fresh window-local ByteTrack tracker is used. ? cv2 frame-
seek
270         accuracy varies by codec (it lands near the requested keyframe then decodes
forward);
271         treat the indices as ±a few frames on real footage ? fine for window-level
features.
272         """
273         self.lazy_load_model()
274         cap = cv2.VideoCapture(video_path)
275         if not cap.isOpened():
276             logger.error(f"Could not open {video_path}")
277             return {"frames": {}, "fps": 30.0, "frame_width": 0.0, "frame_height": 0.0}
278
279         fps = cap.get(cv2.CAP_PROP_FPS)
280         if fps <= 0:
281             fps = 30.0
282         frame_width = cap.get(cv2.CAP_PROP_FRAME_WIDTH)
283         frame_height = cap.get(cv2.CAP_PROP_FRAME_HEIGHT)
284         score_thresh = self.config.get("indexing", {}).get("detector_score_threshold",
0.5)
285
286         step = max(frame_skip, 1)
287         tracker = self.make_tracker(fps / step)
288
289         start_frame = max(0, int(start_frame))
290         frames: Dict[int, List[Tuple[int, Tuple[float, float, float, float]]]] = {}
291         cap.set(cv2.CAP_PROP_POS_FRAMES, start_frame)
292         frame_idx = start_frame
293         while frame_idx <= end_frame:
294             ret, frame = cap.read()
295             if not ret:
296                 break
297             if (frame_idx - start_frame) % step == 0:
298                 frames[frame_idx] = self.detect_and_track(frame, tracker, score_thresh)
299                 frame_idx += 1
300         cap.release()
301         return {"frames": frames, "fps": fps,
302                 "frame_width": float(frame_width), "frame_height": float(frame_height)}

```

15. user (2026-06-13T19:17:43.222Z)

```

1      # docs/
2

```

```

3      Live documentation for the badminton-highlight-indexer.
4
5      ## For new contributors / ramp-up (read these first)
6      - [CODE_MAP.md](CODE_MAP.md) ? Complete navigation map of the codebase and data
contracts.
7      - [OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md) ?
**Authoritative owned-model roadmap** (supersedes the 2026-06-07 strategy docs where they
disagree). [NEXT_STEPS.md](NEXT_STEPS.md) ? the live prioritized "what to do next."
8      - [COMMERCIALIZATION.md](COMMERCIALIZATION.md) ? Living tracker for the commercial
dimension (licensing, codec/output economics, hosting, vendoring, packaging) + PMF risks and the
quality/data-phase priorities. Supersedes the archived `COMMERCIAL_READINESS_REVIEW.md`.
9      - [PMF_MARKET_RESEARCH.md](PMF_MARKET_RESEARCH.md) ? Cited market study: sizing,
customer segments, competitor pricing, and the recommended beachhead (feeds COMMERCIALIZATION
S6).
10     - [PLATFORM_ARCHITECTURE.md](PLATFORM_ARCHITECTURE.md) ? Strategy for the multi-user
"assistant coach" platform: decoupled layers + pluggable storage/compute (BYO), multi-tenant
data model, orchestration, and a phased roadmap.
11     - [DEFERRED_HARDENING_PLAN.md](DEFERRED_HARDENING_PLAN.md) ? Design (pending approval)
for the larger deferred items: packaging, train/eval guardrails, and the async job model (#16,
the first infra slice toward the platform).
12     - [I18N_PLAN.md](I18N_PLAN.md) ? Design (owner-gated) for internationalizing the product
**by design**: Kannada-first, Hindi next, extensible to other Indian languages. Shared
framework-agnostic locale foundation across UI chrome, coach-facing backend text, and the future
NL-query / coach-mode surfaces.
13     - [BACKLOG.md](BACKLOG.md) ? Deferred items with rationale and revisit triggers.
14     - [ROADMAP.md](ROADMAP.md) ? High-level four-phase product direction.
15     - [VIDEO_RECORDING_GUIDELINES.md](VIDEO_RECORDING_GUIDELINES.md) ? Capture best
practices (referenced by reports).
16     - [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md) ? Mental model of the
two-layer detection system.
17     - [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) ? Design (owner-gated) for
fusing the shuttle cue with a person (and later audio/pose) cue in one rally-fusion layer; picks
up ADR-007's parked person cue and the named-but-unbuilt "Gate B."
18     - [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Latest (v8) vendor
labeling guidelines mirrored from the collector's `docs/annotation/`, with the fusion-driven
court-calibration / singles-doubles additions marked `? ADDED (fusion)`.
19     - [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
20     - [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (P1 SHIPPED) for A/B + shadow
+ SxS experimentation on rally-detection variants with **zero production regression** ?
experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift,
rollback = flip a flag). Composes the existing `backend/eval/` machinery (`calibration`,
`regression`, `metrics`); implemented in `backend/eval/experiment.py`.
21     - [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) ? **Living log** of rally-detection
quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-
chronological record; terminal-state snapshots get archived under `archives/quality-
iterations/`.
22
23     ## Practical / ongoing
24     - [TRACKNET_WSL_SETUP.md](TRACKNET_WSL_SETUP.md)
25     - Golden set docs (`GOLDEN_SET`)
26     - Evaluation harness and calibration drivers (in `backend/eval/`)
27
28     ## Historical documents
29     Completed plans, old research, ADRs, and checkpoints live in [`archives/`](archives/).
30
31     See `archives/README.md` for the umbrella policy and `archives/past_projects/README.md`
for classification of small tech debt / starter projects.
32

```

16. assistant (2026-06-13T19:17:46.981Z)

Now let me look at the database substrate (events table, `add\_event`, candidate segments) and the rally_gate net estimation, plus the experiment.py surface.

17. user (2026-06-13T19:17:49.551Z)

```

77:         CREATE TABLE IF NOT EXISTS events (
213:     def add_play_segment(self, video_id: str, start_time: float, end_time: float, server:
Optional[str] = None, receiver: Optional[str] = None, metadata: Optional[Dict[str, Any]] = None)
-> Optional[int]:
227:     def add_event(self, segment_id: int, timestamp: float, event_type: str, player:
Optional[str] = None, details: Optional[Dict[str, Any]] = None) -> bool:
231:         "INSERT INTO events (segment_id, timestamp, type, player, details)
VALUES (?, ?, ?, ?, ?)",
240:     def add_candidate_segment(self, video_id: str, detector: str, start_time: float,
end_time: float,
263:     def link_candidate_to_segment(self, candidate_id: int, segment_id: int) -> bool:
277:     def query_candidate_segments(self, video_id: str, detector: Optional[str] = None,

```

18. user (2026-06-13T19:17:50.009Z)

```

1      """Rally-start gate: suppress between-point false fires using shuttle net-crossings.
2
3      WASB (shuttle-only) over-detects: when a player flicks/tosses the shuttle back to
4      the opponent for service after losing a point, that single trajectory looks like a
5      rally. The discriminator is structural, not duration-based: a real rally has the
6      shuttle crossing the net MULTIPLE times (each shot lands on the other side); a
7      single service feed crosses ~once.
8
9      This module is camera-agnostic and needs NO new model ? it works on the trajectory
10     we already produce (Frame,Visibility,X,Y). It estimates the play axis (the image
11     axis with the larger shuttle spread) and the net position (median shuttle coord on
12     that axis ? the shuttle clusters at the two ends and crosses the net region often,
13     so the median sits near the net), then counts sign changes of (coord - net) across
14     the visible points inside a candidate window, with a deadband to ignore jitter.
15
16     Gate A in the over-fire fix spectrum (B = player-stance state machine, future).
17     Pure functions only ? no I/O, unit-tested; intended as a post-filter on the
18     windows from trajectory_to_action_windows, or to fold into it later.
19     """
20     from __future__ import annotations
21
22     from dataclasses import dataclass
23     from typing import List, Sequence, Tuple
24
25     Interval = Tuple[float, float]
26
27
28     @dataclass
29     class GatedWindow:
30         start: float
31         end: float
32         crossings: int
33         visible_points: int
34         kept: bool
35
36
37     def _spread(vals: Sequence[float]) -> float:
38         """Robust spread = p95 - p5 (ignores a few outlier detections)."""
39         if len(vals) < 2:
40             return 0.0
41         s = sorted(vals)
42         lo = s[max(0, int(0.05 * (len(s) - 1)))]
43         hi = s[min(len(s) - 1, int(0.95 * (len(s) - 1)))]
44         return hi - lo
45
46
47     def _median(vals: Sequence[float]) -> float:
48         if not vals:
49             return 0.0
50         s = sorted(vals)

```

```

51     n = len(s)
52     return s[n // 2] if n % 2 else 0.5 * (s[n // 2 - 1] + s[n // 2])
53
54
55 def estimate_play_axis(points) -> Tuple[str, float, float]:
56     """Return (axis 'x'/'y', net_value, span) from visible points.
57
58     Play axis = the image axis along which the shuttle travels most (larger spread).
59     For a baseline camera that's Y (players top/bottom); for a side camera, X.
60     """
61     xs = [p.x for p in points if p.visible]
62     ys = [p.y for p in points if p.visible]
63     sx, sy = _spread(xs), _spread(ys)
64     if sx >= sy:
65         return "x", _median(xs), sx
66     return "y", _median(ys), sy
67
68
69 def count_net_crossings(window_points, axis: str, net: float, deadband: float) -> int:
70     """Count sign changes of (coord - net) across visible points, ignoring the
71     deadband zone near the net (so jitter at the net doesn't inflate the count)."""
72     crossings = 0
73     last_side = None
74     for p in window_points:
75         if not p.visible:
76             continue
77         v = (p.x if axis == "x" else p.y) - net
78         if abs(v) < deadband:
79             continue
80         side = v > 0
81         if last_side is not None and side != last_side:
82             crossings += 1
83         last_side = side
84     return crossings
85
86
87 def _points_in(points, fps: float, start: float, end: float):
88     return [p for p in points if start * fps <= p.frame <= end * fps]
89
90
91 def gate_windows(points, windows: List[Interval], fps: float,
92                 min_crossings: int = 2, deadband_frac: float = 0.06) ->
List[GatedWindow]:
93     """Annotate each (start,end) window with its shuttle net-crossing count and a
94     keep decision (crossings >= min_crossings). Deadband is a fraction of the play
95     axis span. Net/axis are estimated once over the whole trajectory (stable)."""
96     axis, net, span = estimate_play_axis(points)
97     deadband = deadband_frac * span
98     out: List[GatedWindow] = []
99     for (s, e) in windows:
100         wp = _points_in(points, fps, s, e)
101         c = count_net_crossings(wp, axis, net, deadband)
102         vis = sum(1 for p in wp if p.visible)
103         out.append(GatedWindow(start=s, end=e, crossings=c, visible_points=vis,
104                               kept=c >= min_crossings))
105     return out
106
107
108 def apply_gate(points, windows: List[Interval], fps: float,
109               min_crossings: int = 2, deadband_frac: float = 0.06) -> List[Interval]:
110     """Convenience: return only the windows that pass the crossing gate."""
111     return [(g.start, g.end) for g in
112            gate_windows(points, windows, fps, min_crossings, deadband_frac) if g.kept]
113

```

19. assistant (2026-06-13T19:17:53.211Z)

Let me read the database substrate around events and candidate segments, and the experiment.py public surface.

20. user (2026-06-13T19:17:55.020Z)

```
60         # Play Segments Table (Extensible metadata JSON)
61         cursor.execute("""
62             CREATE TABLE IF NOT EXISTS play_segments (
63                 id INTEGER PRIMARY KEY AUTOINCREMENT,
64                 video_id TEXT NOT NULL,
65                 start_time REAL NOT NULL,
66                 end_time REAL NOT NULL,
67                 duration REAL NOT NULL,
68                 server TEXT,
69                 receiver TEXT,
70                 metadata TEXT,
71                 FOREIGN KEY (video_id) REFERENCES videos(id) ON DELETE CASCADE
72             )
73         """)
74
75         # Events Table
76         cursor.execute("""
77             CREATE TABLE IF NOT EXISTS events (
78                 id INTEGER PRIMARY KEY AUTOINCREMENT,
79                 segment_id INTEGER NOT NULL,
80                 timestamp REAL NOT NULL,
81                 type TEXT NOT NULL,
82                 player TEXT,
83                 details TEXT,
84                 FOREIGN KEY (segment_id) REFERENCES play_segments(id) ON DELETE
85             )
86         """)
87
88         # Versioned migrations live here. PRAGMA user_version is the schema
89         # contract ? bump it for every change and add an ordered `if version < N`
90         # block below. Tests assert the expected version, so accidental drift fails
91         CI.
92         version = cursor.execute("PRAGMA user_version").fetchone()[0]
93
94         if version < 1:
95             # v1: raw candidate detections (pre-confirmation), detector-agnostic.
96             # Common fields are columns; detector-specific metrics go in `stats`
97             JSON,
98             # so a new segmenter reuses this table by setting its own
99             `detector`/`stats`.
100             cursor.execute("""
101                 CREATE TABLE IF NOT EXISTS candidate_segments (
102                     id INTEGER PRIMARY KEY AUTOINCREMENT,
103                     video_id TEXT NOT NULL,
104                     detector TEXT NOT NULL,
105                     chunk_index INTEGER,
106                     start_time REAL NOT NULL,
107                     end_time REAL NOT NULL,
108                     duration REAL NOT NULL,
109                     confidence REAL,
110                     stats TEXT,
111                     detection_params TEXT,
112                     confirmed_segment_id INTEGER,
113                     human_label TEXT,
114                     notes TEXT,
115                     created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
116                     FOREIGN KEY (video_id) REFERENCES videos(id) ON DELETE CASCADE,
```

```

114 FOREIGN KEY (confirmed_segment_id) REFERENCES play_segments(id)
ON DELETE SET NULL
115         )
116     """)
117     cursor.execute("""
118         CREATE INDEX IF NOT EXISTS idx_candidates_video_detector
119         ON candidate_segments(video_id, detector)
120     """)
121     cursor.execute("PRAGMA user_version = 1")
122
123     if version < 2:
124         # v2: async job model (DEFERRED_HARDENING_PLAN #16). One row per
submitted
125         # /api/process request. `status` is the EXECUTION state of the job
126         # (queued|running|done|failed); the pipeline's own verdict (e.g. a video
127         # that failed validation) lives inside `result` JSON as
result["status"].
128         # `result` is the full sync-era response dict (incl. report paths), so
the
129         # observability report is the job's artifact, per the plan.
130         cursor.execute("""
131             CREATE TABLE IF NOT EXISTS jobs (
132                 id TEXT PRIMARY KEY,
133                 video_path TEXT NOT NULL,
134                 video_id TEXT,
135                 segmenter TEXT,
136                 player_pool TEXT,
137                 max_frames INTEGER,
138                 status TEXT NOT NULL DEFAULT 'queued',
139                 progress TEXT,
140                 error TEXT,
141                 result TEXT,
142                 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
143                 started_at TIMESTAMP,
144                 finished_at TIMESTAMP
145             )
146         """)
147         cursor.execute("CREATE INDEX IF NOT EXISTS idx_jobs_status ON
jobs(status)")
148         cursor.execute("PRAGMA user_version = 2")
149         # future: if version < 3: cursor.execute("ALTER TABLE ..."); PRAGMA
user_version = 3
150
151         conn.commit()
152
153     def add_video(self, video_id: str, filepath: str, status: str, validation_results:
List[Dict[str, Any]]) -> bool:
154         """Register or update a video row WITHOUT touching its child segments.
155
156         Uses UPSERT, not INSERT OR REPLACE: REPLACE deletes the conflicting row, which ?
157         with foreign_keys ON ? cascade-deletes every
play_segment/event/candidate_segment.
158         That made a status update (e.g. marking 'failed' on re-validation, or
'processed'
159         before the segmenter runs) silently destroy prior good results. UPSERT mutates
the
160         row in place; clearing segments is now an explicit, deliberate operation
161         (clear_video_data / prune_segments).
162         """
163         try:
164             with self._connect() as conn:
165                 conn.cursor().execute(
166                     "INSERT INTO videos (id, filepath, status, validation_results)
VALUES (?, ?, ?, ?) "
167                     "ON CONFLICT(id) DO UPDATE SET "

```

```

168         "filepath=excluded.filepath, status=excluded.status, "
169         "validation_results=excluded.validation_results",
170         (video_id, filepath, status, json.dumps(validation_results))
171     )
172     conn.commit()
173     return True
174 except Exception as e:
175     logger.error(f"Failed to add video: {e}")
176     return False
177
178 def segment_watermarks(self, video_id: str) -> Tuple[int, int]:
179     """Return (max play_segment id, max candidate_segment id) for a video (0 if
180     none).
181     Captured before a (re)segmentation so the run's new rows (id > watermark) can be
182     told apart from prior rows (id <= watermark), enabling destroy-AFTER-confirm.
183     """
184     with self._connect() as conn:
185         cur = conn.cursor()
186         p = cur.execute("SELECT COALESCE(MAX(id), 0) FROM play_segments WHERE
187         video_id = ?",
188                         (video_id,)).fetchone()[0]
189         c = cur.execute("SELECT COALESCE(MAX(id), 0) FROM candidate_segments WHERE
190         video_id = ?",
191                         (video_id,)).fetchone()[0]

```

21. user (2026-06-13T19:17:55.281Z)

```

205         if play_after is not None:
206             cur.execute("DELETE FROM play_segments WHERE video_id = ? AND id > ?",
207                         (video_id, play_after))
208         if cand_upto is not None:
209             cur.execute("DELETE FROM candidate_segments WHERE video_id = ? AND id <=
210             ?", (video_id, cand_upto))
211         if cand_after is not None:
212             cur.execute("DELETE FROM candidate_segments WHERE video_id = ? AND id >
213             ?", (video_id, cand_after))
214         conn.commit()
215
216 def add_play_segment(self, video_id: str, start_time: float, end_time: float,
217 server: Optional[str] = None, receiver: Optional[str] = None, metadata: Optional[Dict[str, Any]]
218 = None) -> Optional[int]:
219     try:
220         with self._connect() as conn:
221             cursor = conn.cursor()
222             cursor.execute(
223                 "INSERT INTO play_segments (video_id, start_time, end_time,
224                 duration, server, receiver, metadata) VALUES (?, ?, ?, ?, ?, ?, ?)",
225                 (video_id, start_time, end_time, end_time - start_time, server,
226                 receiver, json.dumps(metadata or {})))
227             )
228             conn.commit()
229             return cursor.lastrowid
230     except Exception as e:
231         logger.error(f"Failed to add play segment: {e}")
232         return None
233
234 def add_event(self, segment_id: int, timestamp: float, event_type: str, player:
235 Optional[str] = None, details: Optional[Dict[str, Any]] = None) -> bool:
236     try:
237         with self._connect() as conn:
238             conn.cursor().execute(
239                 "INSERT INTO events (segment_id, timestamp, type, player, details)
240                 VALUES (?, ?, ?, ?, ?)",
241                 (segment_id, timestamp, event_type, player, json.dumps(details or

```



```

{}))
233         )
234         conn.commit()
235         return True
236     except Exception as e:
237         logger.error(f"Failed to add event: {e}")
238         return False
239
240     def add_candidate_segment(self, video_id: str, detector: str, start_time: float,
end_time: float,
241                               chunk_index: Optional[int] = None, confidence:
Optional[float] = None,
242                               stats: Optional[Dict[str, Any]] = None,
243                               detection_params: Optional[Dict[str, Any]] = None) ->
Optional[int]:
244         """Persist one raw (pre-confirmation) detection. `stats`/`detection_params` are
245         detector-specific dicts stored as JSON. Returns the new candidate id, or
246         None."""
247         try:
248             with self._connect() as conn:
249                 cursor = conn.cursor()
250                 cursor.execute(
251                     """INSERT INTO candidate_segments
252                     (video_id, detector, chunk_index, start_time, end_time, duration,
253                     confidence, stats, detection_params)
254                     VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)""",
255                     (video_id, detector, chunk_index, start_time, end_time, end_time -
start_time,
256                     confidence, json.dumps(stats or {}), json.dumps(detection_params or
{}))
257                 )
258                 conn.commit()
259                 return cursor.lastrowid
260         except Exception as e:
261             logger.error(f"Failed to add candidate segment: {e}")
262             return None
263
264     def link_candidate_to_segment(self, candidate_id: int, segment_id: int) -> bool:
265         """Record that a candidate detection was confirmed as a given play_segment."""
266         try:
267             with self._connect() as conn:
268                 conn.cursor().execute(
269                     "UPDATE candidate_segments SET confirmed_segment_id = ? WHERE id =
?",
270                     (segment_id, candidate_id)
271                 )
272                 conn.commit()
273                 return True
274         except Exception as e:
275             logger.error(f"Failed to link candidate {candidate_id} to segment
{segment_id}: {e}")
276             return False
277
278     def query_candidate_segments(self, video_id: str, detector: Optional[str] = None,
only_unlabeled: bool = False) -> List[Dict[str, Any]]:
279         """Fetch candidate detections for the annotation / fine-tuning workflow.
280         `stats` and `detection_params` are deserialized from JSON."""
281         query = "SELECT * FROM candidate_segments WHERE video_id = ?"
282         params: List[Any] = [video_id]
283
284         if detector:
285             query += " AND detector = ?"
286             params.append(detector)
287         if only_unlabeled:
288             query += " AND human_label IS NULL"

```

```

289
290         query += " ORDER BY start_time"
291
292         candidates = []
293         with self._connect() as conn:
294             rows = conn.cursor().execute(query, params).fetchall()
295             for row in rows:
296                 d = dict(row)
297                 d["stats"] = json.loads(d["stats"] or "{}")
298                 d["detection_params"] = json.loads(d["detection_params"] or "{}")
299                 candidates.append(d)
300         return candidates
301
302     def get_video(self, video_id: str) -> Optional[Dict[str, Any]]:
303         with self._connect() as conn:
304             row = conn.cursor().execute("SELECT * FROM videos WHERE id = ?",
305 (video_id,)).fetchone()
306             if row:
307                 d = dict(row)
308                 d["validation_results"] = json.loads(d["validation_results"] or "[]")
309                 return d
310             return None
311
312     def query_play_segments(self, video_id: str, filters: Dict[str, Any]) ->
313 List[Dict[str, Any]]:
314         """
315         Dynamically queries play segments based on filters:
316         - duration_min: float
317         - server: string
318         - receiver: string
319         - event_type: string (e.g. smash, foul)
320         """
321         query = "SELECT r.* FROM play_segments r WHERE r.video_id = ?"
322         params = [video_id]
323
324         if filters.get("duration_min"):
325             query += " AND r.duration >= ?"
326             params.append(filters["duration_min"])

```

22. assistant (2026-06-13T19:17:59.552Z)

I have most of what I need. Let me look at the experiment.py surface and the config models (fusion block / cue_flags) to reference them accurately.

23. user (2026-06-13T19:18:00.484Z)

```

1     """Experiment harness ? A/B + shadow + side-by-side for rally detection.
2
3     This is the thin **variant/experiment layer** that `docs/EXPERIMENT_HARNESS.md`
4     identifies as the one missing piece: ~80% of the machinery (scoring, LOVO, the
5     no-regression gate, the pinnable classifier artifact, the persisted-candidate
6     substrate) already exists ? this module composes it. **No new scoring math lives
7     here**; everything defers to:
8
9     - `backend.eval.metrics`      ? `score`, `match_intervals`, `interval_iou`
10    - `backend.eval.training`      ? `label_candidates` (y by the same IoU matcher)
11    - `backend.eval.classifier`    ? `LogisticRegression` (the pinnable JSON model)
12    - `backend.eval.gemini_refine` ? `merge_overlaps` (the blessed window merge)
13    - `backend.eval.regression`    ? `gt_hash` (label-set fingerprint)
14    - `backend.eval.calibration`   ? `pct_improvement`
15
16    The thesis (`EXPERIMENT_HARNESS.md` §2): separate the **expensive, risky DETECTION**
17    (per-frame signals ? run once on the GPU/WSL box, persisted into
18    `candidate_segments.stats`) from the **cheap, swappable DECISION** (given those
19    signals, where are the rallies). A `Variant` is a declarative re-scoring recipe;

```

```

20 given persisted candidate features it produces `List[Interval]` deterministically,
21 with no GPU and zero production risk. So most experiments are pure offline
22 re-scoring of stored data and never touch the served pipeline.
23
24 Three modes (`EXPERIMENT_HARNESS.md` §5):
25 - `compare()` ? labeled A/B vs golden labels: per-video + LOVO-honest
26   aggregate deltas. Mirrors `distill_local`'s honest train-on-others/predict-
27   held-out loop, but variant-parameterised.
28 - `compare_unlabeled()` ? SxS on NEW footage with no ground truth: where do two
29   variants agree / diverge (A-only, B-only, agreed). The divergences are what a
30   human spot-checks (and what two highlight reels would show side by side).
31 - the promotion gate stays `run_regression.py` + `regression_baseline.json`
32   (exit 0/1/3); **rollback = flip the variant/config, never `git revert`**
33
34 Pure + offline + unit-tested. The only DB-touching helper is
35 `load_video_candidates`, which reads persisted candidates via
36 `Database.query_candidate_segments`.
37 """
38 from __future__ import annotations
39
40 import json
41 import logging
42 import os
43 from dataclasses import dataclass, field
44 from datetime import datetime, timezone
45 from hashlib import sha1
46 from typing import Any, Dict, List, Optional, Sequence
47
48 from backend.eval.calibration import pct_improvement
49 from backend.eval.classifier import LogisticRegression
50 from backend.eval.gemini_refine import merge_overlaps
51 from backend.eval.metrics import Interval, match_intervals, score
52 from backend.eval.regression import gt_hash
53 from backend.eval.training import label_candidates
54
55 logger = logging.getLogger(__name__)
56
57 # Where a comparison run appends one reproducible record. Tracked location (NOT under
58 # the gitignored output/) so the leaderboard survives a clean checkout, mirroring
59 # regression.DEFAULT_BASELINE_PATH.
60 DEFAULT_LEADERBOARD_PATH = os.path.join("eval_baselines", "experiment_leaderboard.json")
61 _METRICS = ("precision", "recall", "f1", "mean_iou")
62
63
64 class ExperimentError(ValueError):
65     """A variant cannot be evaluated as configured (e.g. a learned variant asked to
66     score an unlabeled video with no pinned model, or a gate referencing an absent
67     feature)."""
68
69
70 # ----- Variant
71
72
73 @dataclass
74 class Variant:
75     """A declarative re-scoring recipe ? NOT a code branch (`EXPERIMENT_HARNESS.md` §4).
76
77     A variant turns one video's persisted candidate windows + features into rally
78     intervals. Every knob defaults to the **control** behaviour (no gate, no learned
79     filter), so a variant that merely *carries* a new, disabled cue is byte-identical
80     to control ? the core no-regression guarantee.
81
82     Fields:
83     - segmenter` / config_overrides` / cue_flags` ? informational provenance
84       for the leaderboard and for selecting the served path (the existing

```

```

85         `segmenter_registry` + `IndexingConfig` do the actual selection in production).
86     New cues are recorded here default-OFF.
87     - ``min_crossings`` ? decision-gate **Gate A**: keep only windows whose persisted
88       ``net_crossings`` feature is >= this. 0 = off. This is the eval-only gate from
89       `rally_gate.py` expressed as a re-scoring knob (kills service-feed / held-
shuttle
90       windows ? see `MULTI_SIGNAL_FUSION_PLAN.md` §3.1).
91     - ``min_players`` ? decision-gate **Gate B** (the future person cue): keep windows
92       whose persisted ``players_in_court`` is >= this. 0 = off (no-op until the person
93       cue persists that feature).
94     - ``classifier_model`` ? path to a frozen `LogisticRegression` JSON. When set,
95       `compare()` scores videos with the SHIPPED model as-is (no per-fold training);
96       required for `compare_unlabeled` on a learned variant.
97     - ``feature_names`` ? the persisted features the learned filter consumes. When set
98       WITHOUT ``classifier_model``, `compare()` trains a fresh model per LOVO fold
99       (honest) ? the recipe, not a pinned artifact.
100    - ``threshold`` ? classifier probability cutoff. ``merge_gap`` ? post-filter
101      overlap-merge gap (seconds), the same `gemini_refine.merge_overlaps` default.
102    """
103
104    name: str
105    segmenter: str = "wasb_hybrid"
106    config_overrides: Dict[str, Any] = field(default_factory=dict)
107    cue_flags: Dict[str, bool] = field(default_factory=dict)
108    min_crossings: int = 0
109    min_players: int = 0
110    classifier_model: Optional[str] = None
111    feature_names: Optional[List[str]] = None
112    threshold: float = 0.5
113    merge_gap: float = 1.0
114
115    def is_learned(self) -> bool:
116        """True if this variant applies a learned classifier (pinned model or
recipe)."""
117        return self.classifier_model is not None or bool(self.feature_names)
118
119    def to_dict(self) -> dict:
120        return {
121            "name": self.name,
122            "segmenter": self.segmenter,
123            "config_overrides": self.config_overrides,
124            "cue_flags": self.cue_flags,
125            "min_crossings": self.min_crossings,
126            "min_players": self.min_players,
127            "classifier_model": self.classifier_model,
128            "feature_names": self.feature_names,
129            "threshold": self.threshold,
130            "merge_gap": self.merge_gap,
131        }
132
133    @classmethod
134    def from_dict(cls, d: dict) -> "Variant":
135        known = {f for f in cls.__dataclass_fields__} # type: ignore[attr-defined]
136        return cls(**{k: v for k, v in d.items() if k in known})
137
138    def spec_hash(self) -> str:
139        """Stable 12-char hash of the recipe ? identifies the variant in the leaderboard
140        and makes a result reproducible. The pinned **control** variant always hashes
the
141        same, so a regression is always traceable to a recipe change, never to drift."""
142        blob = json.dumps(self.to_dict(), sort_keys=True, default=str)
143        return sha1(blob.encode()).hexdigest()[:12]
144
145
146    #: The pinned, frozen baseline: candidates as detected, merged, no gate, no learned

```

```

147 #: filter. Always the comparison reference; "rollback" = make this the served variant.
148 CONTROL = Variant(name="control")
149
150
151 # ----- persisted candidates
152
153
154 @dataclass
155 class VideoCandidates:
156     """One video's persisted detection, ready for offline re-scoring.
157
158     ``intervals`` are the candidate windows; ``features`` is column-major
159     (``feature_name -> per-candidate value``, each list aligned to ``intervals``);
160     ``gts`` are golden labels (None for new/unlabeled footage). Build one from the DB
161     with ``load_video_candidates``, or directly in tests.
162     """
163
164     name: str
165     intervals: List[Interval]
166     features: Dict[str, List[float]] = field(default_factory=dict)
167     gts: Optional[List[Interval]] = None
168
169
170 def _feature_column(vc: VideoCandidates, name: str) -> List[float]:
171     """Persisted feature column, defaulting missing/short columns to 0.0 (matches
172     ``training.extract_yolo_features``, which lets a sparse/old row still vectorise)."""
173     col = vc.features.get(name)
174     n = len(vc.intervals)
175     if col is None:
176         logger.warning("variant references feature %r absent from %s ? treating as 0.0",
177                        name, vc.name)
178         return [0.0] * n
179     if len(col) != n:
180         raise ExperimentError(
181             f"feature {name!r} has {len(col)} values but {vc.name} has {n} candidates")
182     return [float(x) for x in col]
183
184
185 def _feature_matrix(vc: VideoCandidates, feature_names: Sequence[str]) ->
List[List[float]]:
186     cols = [_feature_column(vc, name) for name in feature_names]
187     return [list(row) for row in zip(*cols)] if cols else [[] for _ in vc.intervals]
188
189
190 def _gate_mask(variant: Variant, vc: VideoCandidates) -> List[bool]:
191     """Boolean keep-mask from the decision gates (Gate A net-crossings, Gate B
192     players)."""
193     n = len(vc.intervals)
194     mask = [True] * n
195     if variant.min_crossings > 0:
196         col = _feature_column(vc, "net_crossings")
197         mask = [m and (col[i] >= variant.min_crossings) for i, m in enumerate(mask)]
198     if variant.min_players > 0:
199         col = _feature_column(vc, "players_in_court")
200         mask = [m and (col[i] >= variant.min_players) for i, m in enumerate(mask)]
201     return mask
202
203
204 def predict(variant: Variant, vc: VideoCandidates,
205             model: Optional[LogisticRegression] = None) -> List[Interval]:
206     """Variant prediction: gate by persisted features, optionally filter by a learned
207     model, then merge. Pure and deterministic.
208
209     ``model`` (an already-fitted ``LogisticRegression``) is supplied by ``compare`` per LOVO
210     fold; if omitted and the variant pins ``classifier_model``, that frozen model is

```

```

210     loaded. A learned variant with neither a supplied nor a pinned model raises ? there
211     is nothing to score with.
212     """
213     mask = _gate_mask(variant, vc)
214     if variant.feature_names:
215         if model is None:
216             if variant.classifier_model is None:
217                 raise ExperimentError(
218                     f"variant {variant.name!r} is learned but has no model to predict "
219                     f"with (pass a fitted model or set classifier_model)")
220             model = LogisticRegression.load(variant.classifier_model)
221             proba = model.predict_proba(_feature_matrix(vc, variant.feature_names))
222             mask = [m and (float(proba[i]) >= variant.threshold) for i, m in
enumerate(mask)]
223         kept = [iv for iv, keep in zip(vc.intervals, mask) if keep]
224         return merge_overlaps(kept, gap_tol=variant.merge_gap)
225
226
227     # ----- variant scoring
228
229
230     def _train(variant: Variant, videos: Sequence[VideoCandidates], iou: float) ->
LogisticRegression:
231         """Fit the variant's classifier on the pooled candidates of ``videos`` (labels via
the
232         evaluator's own IoU matcher). The honest-LOVO training step from `distill_local`. """
233         X: List[List[float]] = []
234         y: List[int] = []
235         for vc in videos:
236             if vc.gts is None:
237                 raise ExperimentError(f"cannot train: {vc.name} has no golden labels")
238             X.extend(_feature_matrix(vc, variant.feature_names or []))
239             y.extend(label_candidates(vc.intervals, vc.gts, iou))
240         if not X:
241             raise ExperimentError("cannot train a learned variant on zero candidates")
242         return LogisticRegression().fit(X, y, variant.feature_names)
243
244
245     def evaluate_variant(videos: Sequence[VideoCandidates], variant: Variant,
246                         iou: float = 0.5, honest: bool = True) -> Dict[str, Any]:
247         """Score a variant on a labeled corpus. Returns per-video Scores + predictions + a
248         mean aggregate.
249
250         Prediction policy:
251         - **static variant** (no learned filter): predict each video directly ? no
training,
252         so no leakage; per-video scoring is already honest.
253         - **learned, pinned model** (``classifier_model`` set): score every video with the
SHIPPED model. ? optimistic if those videos were in its training set ? use this
254         to
255         report "how the shipped artifact does here", not for the promotion decision.
256         - **learned recipe** (``feature_names``, no pinned model) + ``honest=True``: LOVO
?
257         train on the OTHER videos, predict the held-out one. The number to trust.
258         - learned recipe + ``honest=False``: train on all, predict each (overfit; sanity
only).
259         """
260         for vc in videos:
261             if vc.gts is None:
262                 raise ExperimentError(f"{vc.name} has no golden labels; use
compare_unlabeled")
263
264         per_video: List[Dict[str, Any]] = []
265         predictions: Dict[str, List[Interval]] = {}
266

```

```

267 recipe_lovo = bool(variant.feature_names) and variant.classifier_model is None
268 pinned_model = (LogisticRegression.load(variant.classifier_model)
269                 if variant.classifier_model is not None else None)
270 all_model = None
271 if recipe_lovo and not honest:
272     all_model = _train(variant, videos, iou)
273
274 for i, vc in enumerate(videos):
275     if recipe_lovo and honest:
276         others = [v for j, v in enumerate(videos) if j != i]
277         if not others:
278             raise ExperimentError("LOVO needs >=2 videos; pass honest=False or a
pinned model")
279         model: Optional[LogisticRegression] = _train(variant, others, iou)
280     elif recipe_lovo: # honest=False ? one model over all
videos
281         model = all_model
282     else: # static variant or pinned model
283         model = pinned_model
284     preds = predict(variant, vc, model=model)
285     assert vc.gts is not None # guarded at function top
286     sc = score(preds, vc.gts, iou)
287     per_video.append({"video": vc.name, "num_pred": len(preds), **{m: getattr(sc, m)
for m in _METRICS}})
288     predictions[vc.name] = preds
289
290 aggregate = {m: (sum(p[m] for p in per_video) / len(per_video) if per_video else
0.0)
291              for m in _METRICS}
292 return {
293     "variant": variant.name,
294     "spec_hash": variant.spec_hash(),
295     "is_learned": variant.is_learned(),
296     "honest_lovo": recipe_lovo and honest,
297     "per_video": per_video,
298     "aggregate": aggregate,
299     "predictions": predictions,
300 }
301
302
303 def compare(videos: Sequence[VideoCandidates], variant_a: Variant, variant_b: Variant,
304            iou: float = 0.5, honest: bool = True) -> Dict[str, Any]:
305     """Offline labeled A/B (`EXPERIMENT_HARNESS.md` §5.1). Scores both variants on the
same
306     golden corpus and reports the **B ? A** deltas, per video and in aggregate.
307
308     The deltas use the existing `metrics.score` (per video) +
`calibration.pct_improvement`;
309     learned variants are scored LOVO-honestly. The result is a self-contained record fit
for
310     `record_experiment` ? variant specs + hashes + the label-set fingerprint travel with
it.
311
312     """
313     if len(videos) < 1:
314         raise ExperimentError("compare needs at least one labeled video")
315     eval_a = evaluate_variant(videos, variant_a, iou, honest)
316     eval_b = evaluate_variant(videos, variant_b, iou, honest)
317
318     delta = {m: eval_b["aggregate"][m] - eval_a["aggregate"][m] for m in _METRICS}
319     delta_pct = {m: pct_improvement(eval_a["aggregate"][m], eval_b["aggregate"][m]) for
m in _METRICS}
320
321     by_video_a = {p["video"]: p for p in eval_a["per_video"]}
322     per_video_delta = [
        {"video": p["video"], **{m: p[m] - by_video_a[p["video"]][m] for m in _METRICS}}

```

```

323         for p in eval_b["per_video"]
324     ]
325
326     # One fingerprint over the whole label set ? detects "the deltas were measured
against
327     # different labels" the same way regression.gt_hash guards the no-regression
baseline.
328     all_gts: List[Interval] = [iv for vc in videos for iv in (vc.gts or [])]
329
330     return {
331         "iou": iou,
332         "label_set_hash": gt_hash(all_gts),
333         "num_videos": len(videos),
334         "variant_a": eval_a,
335         "variant_b": eval_b,
336         "delta": delta,
337         "delta_pct": delta_pct,
338         "per_video_delta": per_video_delta,
339     }
340
341
342     # ----- SxS on unlabeled video
343
344
345     def compare_unlabeled(video: VideoCandidates, variant_a: Variant, variant_b: Variant,
346                           agree_iou: float = 0.5) -> Dict[str, Any]:
347         """Side-by-side on NEW footage with no ground truth (`EXPERIMENT_HARNESS.md` §5.2).
348
349         With no labels there is no lift to measure ? instead this surfaces where the two
350         variants **agree vs diverge**, which is exactly what a human reviews (and what two
351         highlight reels would show side by side):
352
353         - ``agreed`` ? windows both variants found (matched at ``agree_iou``);
354         - ``a_only`` ? A fired, B suppressed (B's added precision, if A was over-firing);
355         - ``b_only`` ? B fired, A did not (B's new detections ? real recall or new FPs:
356           the human / a later golden label adjudicates).
357
358         Both variants must be predictable without labels: a learned variant therefore needs
359         a
360         pinned ``classifier_model`` (you cannot train on an unlabeled video). Reuses
361         ``metrics.match_intervals`` ? no new matching logic.
362         """
363         preds_a = predict(variant_a, video)
364         preds_b = predict(variant_b, video)
365         mr = match_intervals(preds_a, preds_b, agree_iou)
366
367         agreed = [{"a": preds_a[m.pred_index], "b": preds_b[m.gt_index], "iou": m.iou}
368                   for m in mr.matches]
369         agree_iou = [m.iou for m in mr.matches]
370         a_only = [preds_a[i] for i in mr.unmatched_pred_indices]
371         b_only = [preds_b[i] for i in mr.unmatched_gt_indices]
372
373         def _dur(ivs: List[Interval]) -> float:
374             return sum(e - s for s, e in ivs)
375
376         return {
377             "video": video.name,
378             "agree_iou": agree_iou,
379             "variant_a": variant_a.name,
380             "variant_b": variant_b.name,
381             "num_a": len(preds_a),
382             "num_b": len(preds_b),
383             "num_agreed": len(agreed),
384             "mean_agree_iou": (sum(agree_iou) / len(agree_iou)) if agree_iou else 0.0,
385             "duration_a": _dur(preds_a),

```



```

385         "duration_b": _dur(preds_b),
386         "agreed": agreed,
387         "a_only": a_only,
388         "b_only": b_only,
389     }
390
391
392 # ----- leaderboard / DB
393
394
395 def record_experiment(result: Dict[str, Any], path: str = DEFAULT_LEADERBOARD_PATH,
396                     timestamp: Optional[str] = None) -> Dict[str, Any]:
397     """Append a compact, reproducible record of a `compare()` result to the JSON
398     leaderboard (`EXPERIMENT_HARNESS.md` §6: experiments are records). Generalises
399     `regression_baseline.json`; deliberately the size of a baseline file, not a service
400     (§9). Returns the row written. ``timestamp`` is injectable for deterministic tests.
401     """
402     row: Dict[str, Any] = {
403         "timestamp": timestamp or datetime.now(timezone.utc).isoformat(),
404         "iou": result.get("iou"),
405         "label_set_hash": result.get("label_set_hash"),
406         "num_videos": result.get("num_videos"),
407         "variant_a": {"name": result["variant_a"]["variant"],
408                      "spec_hash": result["variant_a"]["spec_hash"],
409                      "aggregate": result["variant_a"]["aggregate"]},
410         "variant_b": {"name": result["variant_b"]["variant"],
411                      "spec_hash": result["variant_b"]["spec_hash"],
412                      "aggregate": result["variant_b"]["aggregate"]},
413         "delta": result.get("delta"),
414     }
415     entries: List[Dict[str, Any]] = []
416     if os.path.isfile(path):
417         with open(path, encoding="utf-8") as f:
418             entries = json.load(f)
419     entries.append(row)
420     os.makedirs(os.path.dirname(path) or ".", exist_ok=True)
421     with open(path, "w", encoding="utf-8") as f:
422         json.dump(entries, f, indent=2)
423     logger.info("recorded experiment %s vs %s (?f1=%.3f) -> %s",
424               row["variant_a"]["name"], row["variant_b"]["name"],
425               (result.get("delta") or {}).get("f1", 0.0), path)
426     return row
427
428
429 def load_video_candidates(db: Any, video_id: str, *, detector: Optional[str] = None,
430                          gts: Optional[List[Interval]] = None,
431                          feature_names: Sequence[str] = ()) -> VideoCandidates:
432     """Bridge persisted detection ? `VideoCandidates` via
433     `Database.query_candidate_segments`.
434     ``feature_names`` selects which persisted features to lift out of each row's
435     ``stats``
436     JSON (``"duration"`` is derived from the window bounds). This is the only DB-
437     touching
438     function; the comparison itself is pure and offline. Exercised owner-side on a real
439     DB,
440     but works against any `Database` (including an in-memory one in tests).
441     """
442     rows = db.query_candidate_segments(video_id, detector=detector)
443     intervals: List[Interval] = [(float(r["start_time"]), float(r["end_time"])) for r in
444     rows]
445     features: Dict[str, List[float]] = {}
446     for name in feature_names:
447         if name == "duration":
448             features[name] = [float(r["end_time"]) - float(r["start_time"]) for r in

```

```

rows]
445         else:
446             features[name] = [float((r.get("stats") or {}).get(name, 0.0)) for r in
rows]
447         return VideoCandidates(name=video_id, intervals=intervals, features=features,
gts=gts)
448

```

24. user (2026-06-13T19:18:00.586Z)

```

1     """Pydantic models for the application configuration.
2
3     This module defines the canonical, validated, typed configuration for the
4     Badminton Highlight Indexer. Built-in defaults live here (single source of truth).
5
6     The models support graceful migration from old config.json files via extra="allow".
7     """
8
9     from __future__ import annotations
10
11     from typing import Any, List, Literal, Optional
12
13     from pydantic import BaseModel, Field, field_validator, model_validator
14
15     Sport = Literal["badminton", "tennis", "pickleball", "table_tennis", "cricket"]
16
17
18     class ValidationRelevanceConfig(BaseModel):
19         court_green_lower: list[int] = Field(default=[35, 40, 40])
20         court_green_upper: list[int] = Field(default=[85, 255, 255])
21         court_pixels_min_ratio: float = Field(default=0.05, ge=0.0, le=1.0)
22
23
24     class ValidationConfig(BaseModel):
25         min_resolution_width: int = 1280
26         min_resolution_height: int = 720
27         min_fps: float = 29.9
28         min_blur_threshold: float = 20.0
29         max_scene_cuts_per_minute: float = 0.5
30         relevance: ValidationRelevanceConfig =
Field(default_factory=ValidationRelevanceConfig)
31
32
33     class TrackNetConfig(BaseModel):
34         wsl_distro: str = "Ubuntu"
35         repo_dir: str = "~/models/TrackNetV4"
36         conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
37         conda_env: str = "TrackNetV4"
38         weights_path: str = ""
39         queue_length: int = 5
40         stage_in_wsl: bool = True
41         wsl_stage_dir: str = "~/clips"
42         timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
43         velocity_threshold: float = 0.02
44         # Cache hygiene: after a SUCCESSFUL run the staged video copy (and, for WASB, the
45         # decoded frame cache) are deleted automatically ? they are pure intermediates,
46         # regenerable from the source, and the dominant disk consumer (~GBs/video). Set
47         # true only for debugging. On FAILURE they are always kept so a re-run can resume.
48         keep_frames: bool = False
49         # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB.
50         # 0 = disabled.
51         min_free_gb: float = 0.0
52
53
54     class WasbIndexConfig(BaseModel):

```

```

55     """Typed config for the live WASB shuttle substrate (indexing.wasb).
56
57     Mirrors backend/pipeline/detectors/wasb_runner.py::WasbConfig.from_indexing_cfg so
58     these knobs actually take effect ? previously this whole block was silently dropped
59     because nested config sections defaulted to extra='ignore'.
60     """
61     wsl_distro: str = "Ubuntu"
62     repo_dir: str = "~/models/WASB-SBDT"
63     conda_sh: str = "~/miniconda3/etc/profile.d/conda.sh"
64     conda_env: str = "wasb"
65     weights_path: str = "~/models/WASB-
SBDT/pretrained_weights/wasb_badminton_best.pth.tar"
66     sport: str = "badminton"
67     wsl_stage_dir: str = "~/clips"
68     timeout_sec: int = 1800 # 30 min; kills a hung WSL/GPU call (0 = no timeout)
69     velocity_threshold: float = 0.02
70     # Cache hygiene: after a SUCCESSFUL run the staged video copy + decoded frame cache
71     # (the ~GBs/video disk hog) are deleted automatically ? pure intermediates,
regenerable
72     # from the source. Set true only for debugging. On FAILURE they are kept so a re-run
resumes.
73     keep_frames: bool = False
74     # Warn before a run if WSL free space (at wsl_stage_dir) is below this many GB. 0 =
disabled.
75     min_free_gb: float = 0.0
76
77
78     class FusionConfig(BaseModel):
79         """Config for the multi-signal `fusion_hybrid` segmenter (MULTI_SIGNAL_FUSION_PLAN
P2).
80
81         Every default reproduces control behaviour: the fusion segmenter persists the
82         shuttle net-crossing count (Gate A signal) and ? only when ``cue_flags['person']``
is
83         true ? the person-cue features, but it does NOT gate the served handoff unless a
84         threshold is raised (``min_crossings``/``min_players`` default 0 = OFF). The court
mask
85         is optional: ``court_polygon=None`` ? Gate B counts every detection (player-count-
only);
86         supplied ? off-court persons (spectators / adjacent courts) are excluded.
87         """
88         # Which trajectory substrate seeds the windows (shuttle stays primary).
89         substrate: Literal["wasb", "tracknet"] = "wasb"
90         # Windowing velocity threshold for the fusion family (the engine reads
91         # indexing.<family>.velocity_threshold). Default matches the substrates' 0.02; set
92         # equal to indexing.wasb.velocity_threshold to A/B fusion against a tuned wasb run.
93         velocity_threshold: float = 0.02
94         # Per-cue enable flags, e.g. {"person": true}. Person cue is OFF unless explicitly
95         # enabled ? so the default path never imports torch / touches the GPU.
96         cue_flags: dict[str, bool] = Field(default_factory=dict)
97         # Served Gate A: drop windows whose shuttle net-crossings < this from the AI
handoff.
98         # 0 = OFF (no gating; persisted candidates are always the full superset for offline
re-scoring).
99         min_crossings: int = Field(default=0, ge=0)
100         # Served Gate B: when the person cue is on, drop windows with median in-court
players
101         # < this. 0 = OFF.
102         min_players: int = Field(default=0, ge=0)
103         # Optional court mask as normalised 0-1 [[x, y], ...] polygon. None = no mask.
104         court_polygon: Optional[list[list[float]]] = None
105         # Net line for the person two-sided-motion split (person features only ? the shuttle
106         # net-crossing gate keeps rally_gate's camera-agnostic auto-estimate).
107         net_axis: Literal["x", "y"] = "y"
108         net_position: float = Field(default=0.5, ge=0.0, le=1.0)

```

```

109
110
111 class IndexingConfig(BaseModel):
112     default_segmenter: str = "yolo_hybrid"
113     use_gpu: bool = True
114     motion_threshold: float = 0.08
115     min_segment_duration: float = 3.0
116     dead_time_merge_gap: float = 2.0
117     chunk_duration: float = 90.0
118     chunk_overlap: float = 10.0
119     detector_model: str = "fasterrcnn_resnet50_fpn"
120     detector_score_threshold: float = 0.5
121     yolo_velocity_threshold: float = 0.15
122     yolo_frame_skip: int = 3
123     yolo_tracker: str = "bytetrack.yaml"
124     yolo_prefetch_workers: int = 2
125     min_action_window_duration: float = 2.0
126     log_yolo_candidates: bool = True
127     store_yolo_candidates: bool = True
128     ai_handoff_padding: float = 2.0
129     ai_max_workers: int = 5
130     # Width (px) the gemini segmenter re-encodes chunks to before upload. Stream-copied
131     # chunks of a high-bitrate (4K) source blow past the provider upload cap
132     # (backend/providers/gemini.py::MAX_UPLOAD_MB) and are refused ? observed live as
133     # "0 segments / success". Re-encoding also makes chunk boundaries frame-accurate
134     # (stream copy cuts at the nearest keyframe). 0 = legacy stream-copy at source res.
135     # Matches gemini_refine's --clip-scale-width default.
136     ai_chunk_scale_width: int = 1280
137     # Optional debug knob: trim every video to the first N seconds before processing.
138     debug_duration: Optional[float] = None
139
140     tracknet: TrackNetConfig = Field(default_factory=TrackNetConfig)
141     wasb: WasbIndexConfig = Field(default_factory=WasbIndexConfig)
142     fusion: FusionConfig = Field(default_factory=FusionConfig)
143
144     @field_validator("default_segmenter")
145     @classmethod
146     def segmenter_must_be_registered(cls, v: str) -> str:
147         # Registry check happens at runtime in main/segmenter selection.
148         # Here we just allow any string for forward compatibility.
149         return v
150
151     @model_validator(mode="after")
152     def _chunk_step_must_be_positive(self) -> "IndexingConfig":
153         # The chunked segmenters advance by (chunk_duration - chunk_overlap); a non-
154         # positive
155         # step makes `while t < duration: t += step` loop forever, growing memory before
156         # any
157         # GPU work. Reject the bad config at load time instead.
158         if self.chunk_overlap >= self.chunk_duration:
159             raise ValueError(
160                 f"indexing.chunk_overlap ({self.chunk_overlap}) must be < chunk_duration
161                 "
162                 f"({self.chunk_duration}); otherwise chunking never advances."
163             )
164         return self
165
166 class OutputConfig(BaseModel):
167     default_highlights_name: str = "highlights.mp4"
168     db_name: str = "sports_indexer.db"
169     # Directory where the DB, highlight reels, and processed clips are written.
170     # The CLI's --output-dir overrides this at startup (see backend/main.py::main).
171     output_dir: str = "output"

```

```

171
172 class WatermarkConfig(BaseModel):
173     """Branding overlay burned into each highlight clip during the per-clip re-encode
174     (backend/pipeline/stitcher/compiler.py). **On by default** ? set `enabled: false`
175     (under `stitching.watermark` in config.json) to turn it off. The text is fully
176     configurable. The concat stage stays stream-copy (-c copy)."""
177     enabled: bool = True
178     text: str = "Generated by Khelsutra.guru"
179     logo_path: str = "" # optional transparent PNG overlay; "" = no logo
180     font_path: str = "" # optional .ttf; "" = use the fontconfig `font` family
below
181     font: str = "Sans" # fontconfig family used when font_path is empty
182     font_size_ratio: float = Field(default=0.035, gt=0.0, le=0.5) # text height as a
fraction of frame height
183     opacity: float = Field(default=0.85, ge=0.0, le=1.0)
184     box: bool = True # faint backing box behind the text for legibility
185     margin: int = Field(default=24, ge=0) # px inset from the chosen corner
186     position: Literal["bottom-right", "bottom-left", "top-right", "top-left"] = "bottom-
right"
187
188
189 class StitchingConfig(BaseModel):
190     begin_padding: float = 0.5
191     end_padding: float = 0.5
192     use_cache: bool = False
193     min_shot_count: int = 1
194     watermark: WatermarkConfig = Field(default_factory=WatermarkConfig)
195
196
197 class SecurityConfig(BaseModel):
198     # When set, /api/process_video_path must resolve under this directory (hosted/tenant
199     # mode). Default None preserves the single-user local 'any local file' workflow.
200     allowed_video_root: Optional[str] = None
201
202     # --- Chunked upload (B2, PR-2) ---
203     # Landing zone for browser uploads (assembled from parts). Per-user subdirectories
204     # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
205     # contain upload_dir or /api/process will reject the uploaded paths.
206     upload_dir: str = "output/uploads"
207     # Whole-file cap (default 4 GB) and per-part cap. Parts stay under ~100 MB because
208     # free-tier edge proxies cap request bodies there (HOSTING_PLAN §2).
209     max_upload_mb: int = 4096
210     max_part_mb: int = 95
211     allowed_upload_extensions: List[str] = ["mp4", "mov"]
212
213
214 class AppConfig(BaseModel):
215     """Root configuration object.
216
217     All runtime configuration should be accessed via an instance of this class
218     (or the loader) rather than raw dict lookups.
219     """
220
221     sport: Sport = "badminton"
222     ai_provider: str = "gemini"
223     # Auto-tracking alias: 2.5-flash is legacy now that 3.5-flash is the stable
flagship;
224     # the alias never deprecates underneath us. Eval reports record the run date
alongside.
225     ai_model: str = "gemini-flash-latest"
226
227     validation: ValidationConfig = Field(default_factory=ValidationConfig)
228     indexing: IndexingConfig = Field(default_factory=IndexingConfig)
229     output: OutputConfig = Field(default_factory=OutputConfig)
230     stitching: StitchingConfig = Field(default_factory=StitchingConfig)

```

```

231     security: SecurityConfig = Field(default_factory=SecurityConfig)
232
233     # Allow extra keys during transition so existing config.json files do not break.
234     model_config = {
235         "extra": "allow",
236         "validate_assignment": True,
237     }
238
239     # Temporary migration shim for legacy dict-style access (config.get("section",
dict    {})).get(...)).
240     # Will be removed once all call sites (validators, segmenters, etc.) move to typed
241     # attribute access / the typed config. A nested *section* is returned as a plain
dict
242     # via model_dump, so legacy chained access keeps working ? returning the Pydantic
243     # section model instead breaks it (a model has no .get), which bit every validator
on
244     # real runs. Scalar leaves are returned as-is.
245     def get(self, key: str, default: Any = None) -> Any:
246         _MISSING = object()
247         val: Any = _MISSING
248         if key in type(self).model_fields or hasattr(self, key):
249             val = getattr(self, key, _MISSING)
250         if val is _MISSING:
251             # Legacy convenience: find a leaf key inside any nested section. Sections
are
252             # discovered dynamically from the model fields, so a future top-level
section
253             # (e.g. `audio`, `jobs`) is covered automatically without editing a
hardcoded list.
254             for fname in type(self).model_fields:
255                 section = getattr(self, fname, None)
256                 if isinstance(section, BaseModel) and hasattr(section, key):
257                     val = getattr(section, key)
258                     break
259             if val is _MISSING:
260                 return default
261             if isinstance(val, BaseModel):
262                 return val.model_dump(mode="json")
263             return val
264
265     def __getitem__(self, key: str) -> Any:
266         val = self.get(key)
267         if val is None and key not in type(self).model_fields:
268             # Re-raise like a real dict for unknown top-level
269             raise KeyError(key)
270         return val
271
272
273     # Convenience type alias for callers that want to annotate
274     Config = AppConfig
275

```

25. assistant (2026-06-13T19:18:04.506Z)

I have a deep understanding now. Let me check the ADR-007 reference and the segmenter registry pattern to cite them accurately.

26. user (2026-06-13T19:18:06.993Z)

```

172- player-motion features; medium-high ceiling; moderate effort. *Recommended primary.*
173- **WASB shuttle-motion ? rally detector** ? highest ceiling (the shuttle defines play),
174- fully offline but heaviest (needs the WSL2 WASB runner, ADR-001); can't train WASB
175: itself (no xy labels), only a rally layer on its output.
176-
177-### Decision

```

```

178-
--
292-
293-
294-
295:## ADR-007: Audio "thwack" hit-detection (fused with WASB) is the next rally-detection cue;
pose/stance parked
296-
297-Date: 2026-06-07
298-Status: Accepted (direction); experiments gated on golden eval labels
--
345-separate signal fused in our code**? The answer today is the latter ? and understanding
*why* sets up the
346-long-term route.
347-
348-Why fusion lives in our rally layer, not inside WASB (Option A ? now).WASB is a
frozen, vision-only
349-shuttle detector (image frames ? per-frame shuttle `(x, y, visibility)`; MIT weights we
*run*, not train).
350-It has no audio pathway and it does not decide rallies at all ? rally detection is our own
layer** on top of
351-its trajectory (`trajectory_to_action_windows` ? gate ? the distilled classifier). Making
WASB consume audio

```

27. user (2026-06-13T19:18:07.334Z)

```

1      from abc import ABC, abstractmethod
2      from typing import List, Dict, Any, Type, Optional
3      from backend.infrastructure.database import Database
4      from backend.results import Result # new standardized contract (Result[Value] ==
Success[Value] | Failure)
5
6      class BasePlaySegmenter(ABC):
7          def __init__(self, db: Database, config: Dict[str, Any]):
8              self.db = db
9              self.config = config
10
11          @property
12          @abstractmethod
13          def name(self) -> str:
14              """Unique identifier for this segmenter."""
15              pass
16
17          @property
18          @abstractmethod
19          def description(self) -> str:
20              """Human-readable explanation of this segmenter."""
21              pass
22
23          @abstractmethod
24          def process_video(self, video_id: str, video_path: str, player_pool: List[str] =
None, max_frames: int = None) -> "Result[List[Dict[str, Any]]]":
25              """
26              Processes a video to detect play segments, populates SQLite database,
27              and returns a Result.
28
29              Use isinstance(result, Failure) (or result.is_success()) to distinguish
30              "no segments found" from "something went wrong" (missing key, healthcheck
31              failure, etc.). This is the start of the standardized error/result contract.
32              """
33              pass
34
35      class SegmenterRegistry:
36          def __init__(self):
37              self._segmenters: Dict[str, Type[BasePlaySegmenter]] = {}

```

```

38
39     def register(self, segmenter_cls: Type[BasePlaySegmenter]):
40         """Registers a new segmenter class under its name property."""
41         try:
42             temp_inst = segmenter_cls(None, {})
43             name = temp_inst.name
44             except Exception as e:
45                 raise ValueError(f"Failed to resolve name property from segmenter class
{segmenter_cls.__name__}: {e}") from e
46             self._segmenters[name] = segmenter_cls
47
48     def get(self, name: str) -> Optional[Type[BasePlaySegmenter]]:
49         return self._segmenters.get(name)
50
51     def list_available(self) -> Dict[str, str]:
52         """Returns a dict of registered segmenter names and their descriptions."""
53         res = {}
54         for name, cls in self._segmenters.items():
55             try:
56                 temp_inst = cls(None, {})
57                 desc = temp_inst.description
58             except Exception:
59                 desc = "No description available."
60             res[name] = desc
61         return res
62
63     segmenter_registry = SegmenterRegistry()
64

```

28. assistant (2026-06-13T19:20:31.125Z)

I now have a complete picture of both seams' anchor points: the `\_enrich\_candidate\_stats` / `\_select\_for\_handoff` hooks, the `experiment.py` `Variant`/`predict`/`compare` surface (note `\_gate\_mask` already treats `net\_crossings`/`players\_in\_court` as feature columns ? the exact "signals as columns, not branches" pattern), the `candidate\_segments.stats` + `events` substrate, and the `PRAGMA user\_version` migration ladder. Let me set up the branch and write the doc.

29. user (2026-06-13T19:20:32.077Z)

Chapter marked: "Analyzers & reconcilers design doc". Continue your current work.

30. user (2026-06-13T19:20:33.285Z)

Switched to a new branch 'docs/analyzers-and-reconcilers'
docs/analyzers-and-reconcilers

31. user (2026-06-13T19:22:21.732Z)

File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\docs\ANALYZERS_AND_RECONCILERS.md (file state is current in your context ? no need to Read it back)

32. assistant (2026-06-13T19:20:36.836Z)

Now I'll write the design doc.

33. assistant (2026-06-13T19:22:21.699Z)

Analizers & reconcilers ? two extensibility seams for rally-detection quality

****Status:**** ****DESIGN** ? owner-gated, NOT started. ****** A peer to [EXPERIMENT_HARNESS.md] (EXPERIMENT_HARNESS.md) and


```
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md).
```

This doc specs two **seams** ? a forward one (injectable window analyzers) and a backward one (a reconciliation pass) ? that let us keep piling on clever rally signals ***without the pipeline becoming a pile of special cases.*** Docs-only; no implementation code. Each phase below is a separate, owner-gated PR.

```
> **Companion docs:** the fusion seam these plug into is
[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md)
> (`RallyCue` = swappable producer; the `_enrich_candidate_stats` hook; the scale/translation-
invariant
> feature discipline, §3.2); the no-regression A/B + LOVO measurement these ride is
> [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) (`Variant`, `compare`, `compare_unlabeled`);
the scorer
> is the distilled logistic regression (**ADR-004**, `backend/eval/classifier.py`); the modular
rally
> layer is **ADR-007** ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the
2-layer
> mental model is [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md); the live quality
log is
> [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md).
```

1. The spine ? NO special-casing in the decision path

This is the one principle everything else serves:

```
> **Every analyzer and reconciler emits a SIGNAL carrying a value + a confidence/presence. The
SCORING
> MODEL consumes those signals and *learns to weight them*. We do NOT add hand-coded `if/else`
branches
> into the decision.**
```

Analyzers propose; the scorer arbitrates; the experiment harness + golden set decide what gets promoted. That separation is what lets us add a net-hit service fault, a let-cord, a double-bounce, a shuttle-out cue ? each as *"register a producer + a feature column,"* never *"add a branch."* A signal that doesn't earn its keep stays at weight ? 0; it costs nothing because it is an unused column, not a live *`if`*.

The pattern already exists in miniature. *`experiment.predict()`* (*`backend/eval/experiment.py:203`*) treats *`net\_crossings`* and *`players\_in\_court`* as ***feature columns*** ? *`\_gate\_mask`* is the trivial scorer (threshold a column), *`LogisticRegression`* is the richer learned scorer (weight all columns), and a future scorer is richer still. The decision path is **signals ? scorer**, with ***zero*** *`if`* *service_fault*: drop`. Analyzers feed that path more columns; reconcilers run **outside** it as an audited post-pass.

...

```
???????????? DETECTION (expensive, GPU/WSL, run ONCE) ??????????
frames ?? WASB (frozen, MIT) ?? shuttle (x,y,v,vis)/frame ??
frames ?? PersonDetector (ours)?? boxes + tracks/frame ?????? per-frame-aligned
audio ?? hit detector (ours) ?? hit events ?????????????????? signals (+ net estimate,
???????????????????????????????????????????????????????? window bounds)
?
shuttle-seeded candidate windows (FusionSegmenter, P2)
?
???????????? ANALYZERS ? forward seam, run PER WINDOW (flag-gated, default OFF) ??????????
```

```

? WindowAnalyzer registry: service-fault · double-bounce · let-cord · shuttle-out · ? ?
? each ? AnalyzerResult{ signals(value+confidence), events, optional boundary hints } ?
????????????????????????????????????????????????????????????????????????????????
? signals merged into candidate_segments.stats
?
DECISION / SCORING ? signals ? learned scorer (NO branches)
_gate_mask (trivial) + LogisticRegression (ADR-004) + harness-pinned model
? predicted rally segments
?
???????????? RECONCILERS ? backward seam, "linter for rally decisions" ?????????????????
? ConsistencyRule registry (fixed precedence, idempotent): decision-vs-evidence conflicts ?
? ? Issue + corrected set. REPORT-FIRST: report always; apply only under flag / review. ?
????????????????????????????????????????????????????????????????????????????????
?
candidate_segments.stats · events table · corrections record · highlights
?
EXPERIMENT HARNESS (compare / LOVO / golden set) ? gates EVERY promotion
...
---
```

2. Seam 1 ? injectable window analyzers (forward)

A registry of pluggable `WindowAnalyzer`s ? mirroring `segmenter\_registry`
(`backend/pipeline/segmenters/base.py:35`) ? that run over a candidate window's **per-frame-aligned**
signals, all already produced^{**}: the WASB shuttle trajectory (`.frame/.visible/.x/.y` +
velocity), the
person boxes per frame (`PersonDetector.detections\_per\_frame`, `person\_detector.py:256`), the
net line
(`rally\_gate.estimate\_play\_axis`, camera-agnostic), and the window bounds. The analyzer is
^{**pure} ? no
I/O, no torch/cv2 ? mirroring `fusion\_features.py`'s discipline, so its math is unit-testable in
the
GPU-free dev container.

2.1 Contracts (JSON-serializable; proposed `backend/pipeline/analyzers/`)

```

```python
@dataclass(frozen=True)
class Signal:
 """One scalar cue + how sure the analyzer is. Persisted as two stats columns:
 `<analyzer>__<key>` (value) and `<analyzer>__<key>_confidence` (presence, [0,1])."""
 value: float
 confidence: float = 1.0 # 1.0 = deterministic/certain

@dataclass(frozen=True)
class AnalyzerEvent: # ? the `events` table, on confirmation ($4)
 type: str # "net_hit", "double_bounce", "let_cord", ?
 timestamp: float # seconds, source-absolute
 confidence: float = 1.0
 player: Optional[str] = None
 details: Dict[str, Any] = field(default_factory=dict)

@dataclass(frozen=True)
class BoundaryHint: # OPTIONAL deterministic boundary proposal
 kind: Literal["start", "end"]
 timestamp: float
 confidence: float
 reason: str # audit string ? never silently applied

@dataclass(frozen=True)
class AnalyzerResult:
 signals: Dict[str, Signal] = field(default_factory=dict)
```

```

events: List[AnalyzerEvent] = field(default_factory=list)
boundary_hints: List[BoundaryHint] = field(default_factory=list)

@dataclass
class WindowSignals:
 # the read-only, pre-produced bundle (no I/O inside)
 start: float; end: float; fps: float
 frame_width: float; frame_height: float
 shuttle: List["TrajectoryPoint"] # in-window points (.frame/.visible/.x/.y)
 persons: Dict[int, List[Tuple[int, BBox]]] # frame_idx ? [(track_id, box)]; {} when person
 cue OFF
 net: Tuple[str, float] # (axis, position) ? auto-estimate or config
 base_stats: Dict[str, float] # motion stats already computed for the window

class WindowAnalyzer(ABC):
 name: str # registry key + stats-column prefix
 version: str # bump on math change ? provenance in stats/leaderboard
 feature_keys: List[str] # the columns this emits (declared for the harness)
 event_types: List[str]
 @abstractmethod
 def analyze(self, w: WindowSignals) -> AnalyzerResult: ...
...

```

A `Signal` is **value + confidence**, never a verdict. The three roles an `AnalyzerResult` can fill:

```

| Role | Sink | Consumed by |
|---|---|---|
| (a) features (with confidence) | `candidate_segments.stats` as `<analyzer>__<key>` (+ `_confidence`) | the scoring model and the harness, as feature columns (§3) |
| (b) events | the `events` table (buffered on the candidate, flushed on confirmation, §4) | history / highlights / audit |
| (c) boundary hints (optional) | stashed in `stats`; never auto-applied | the reconciler (§5) / a future boundary-aware scorer |

```

## 2.2 Worked example ? the SERVICE-FAULT analyzer

The owner's held-shuttle bug has a sibling: a **service fault** where the serve hits the net and the shuttle comes to rest in the net region. Pure trajectory, GPU-free, no person cue needed:

```

- Input: the window's shuttle trajectory + the `rally_gate` net estimate.
- Logic: the shuttle enters the net region (|coord ? net| < deadband) and its velocity
 decays to ? 0 there (comes to rest) ? high-probability net-hit service fault. Confidence scales with
 how cleanly the halt sits inside the net band.
- Emits: a `service_fault` signal (value = presence, confidence = halt cleanliness) + a `net_hit`
 event at the halt timestamp + a rally-END `BoundaryHint` at that timestamp.

```

Crucially: the analyzer does **not** drop the window. It **proposes** `service\_fault`; the scorer sees `service\_fault` + `service\_fault\_confidence` as two more columns and **learns** whether/how much to down-weight. The boundary hint is an audited proposal the reconciler may act on ? never a branch.

## 2.3 Injection point ? generalize `\_enrich\_candidate\_stats`

The FusionSegmenter hook (`fusion\_hybrid.py:70`) already turns "compute fixed features" into per-window enrichment. We generalize it to **"run each enabled analyzer"** ? additively, so the existing `net\_crossings`/person paths are unchanged:

```

```python

```

```

def _enrich_candidate_stats(self, cw, points, fps, frame_width, video_path, idx_cfg):
    stats = super()._enrich_candidate_stats(...) # base motion keys
    stats["net_crossings"] = ... # Gate-A signal (kept)
    # ? person features when the person cue is on (kept) ?
    w = self._window_signals(cw, points, fps, frame_width, video_path, idx_cfg)
    for analyzer in analyzer_registry.enabled(idx_cfg): # NONE enabled ? byte-identical to
control
        res = analyzer.analyze(w)
        for key, sig in res.signals.items():
            stats[f"{analyzer.name}__{key}"] = sig.value
            stats[f"{analyzer.name}__{key}_confidence"] = sig.confidence
        stats.setdefault("_events", []).extend(e.__dict__ for e in res.events) # flushed
in Phase 4
        stats.setdefault("_boundary_hints", []).extend(h.__dict__ for h in res.boundary_hints)
    return stats
...

```

Flag-gated, ****default OFF**** via the existing ``indexing.fusion.cue_flags`` (a sibling ``analyzer_flags`` map keyed by ``analyzer.name``). With every analyzer OFF, the loop is empty ? candidate rows are byte-identical to the substrate ? the no-regression invariant holds (§6). ``net_crossings`` and the person features are the ****prototype analyzers**** the seam retrofits first (proving it with zero behavior change); every new cue after that is a pure add.

3. How signals reach the scorer ? columns, not branches

The persisted ``stats`` columns are exactly what the harness already lifts. ``load_video_candidates(..., feature_names=[...])`` (``experiment.py:429``) reads each name out of the row's ``stats`` JSON into a feature column; ``predict()`` (``experiment.py:203``) gates/scores on those columns and merges. So a new analyzer signal becomes a scorer input by ****adding its column name to a ``Variant``'s ``feature_names``**** ? no new scoring code, no branch:

```
```python
```

#### **A learned Variant that lets the scorer weight the new service-fault signal:**

```

fusion_plus = Variant(name="fusion+service_fault",
 feature_names=[*SHUTTLE_FEATURES, "person__players_in_court",
 "service_fault__service_fault",
 "service_fault__service_fault_confidence"])
compare(golden_videos, CONTROL, fusion_plus) # LOVO-honest B ? A on the 6 golden videos
...

```

The model **\*\*learns the weights\*\*** ? e.g. down-weight a shuttle-motion window with a high-confidence ``service_fault``, up-weight a window with ?? in-court players and multiple net-crossings. **\*\*What the classifier cannot see, it cannot special-case;\*\*** what it can see, it weights from data. The confidence column lets it discount an uncertain analyzer instead of trusting a binary verdict ? the small-N-safe way to fold in a deterministic-ish cue. (Same logistic regression as ADR-004; we add columns, never a net.)

---

### 4. Persistence

```

- Features ? `candidate_segments.stats`** (`database.py:240`). Namespaced `<analyzer>__<key>`
 (+ `_confidence`). Persisted as a shadow even when the served decision ignores them, so
 any future
 `Variant` re-scores offline (EXPERIMENT_HARNESS §5.2). `version`/`name` recorded for
 provenance.
- Events ? the `events` table** (`database.py:227`, `add_event(segment_id, ts, type, player,
 details)`).
 Its FK is to `play_segments`, which exist only after the AI handoff confirms a candidate.
 So
 analyzer events are buffered on the candidate (`stats["_events"]`) and flushed to
 `events` in
 Phase 4, when `_candidate_id` ? `segment_id` is linked (`trajectory_hybrid.py:283`). Candidate-
 level
 discoveries survive for offline re-scoring; the `events` table stays "events on confirmed
 segments."
- Corrections record** (§5) ? report-first means corrections are recorded, never silent.
 Phase one:
 a JSON sidecar shaped like `experiment_leaderboard.json` (one row per issue: segment, kind,
 evidence,
 before?after, reason, timestamp). When auto-apply is promoted, add a `corrections` table via
 the
 `PRAGMA user_version` migration ladder (`database.py:88`, `if version < 3:`) so applied
 corrections are
 queryable and reversible.

```

---

## 5. Seam 2 ? the reconciliation pass (backward): a linter for rally decisions

Given **predicted segments + the persisted per-frame/per-window signals**, detect where the  
**DECISION**  
 contradicts the EVIDENCE\*\* and propose corrections. This is the backward dual of the analyzers:  
 analyzers  
 add evidence *before* the decision; reconcilers audit the decision *against* evidence *after*  
 it.

### 5.1 Contracts (proposed `backend/pipeline/reconcilers/`)

```

```python
@dataclass(frozen=True)
class Issue:
    segment_index: int
    kind: str                                # rule id, e.g. "trailing_dead_time"
    evidence: Dict[str, Any]                 # the signal values that triggered it (auditable)
    before: Interval
    after: Optional[List[Interval]]          # None = drop · [iv] = corrected · [iv1, iv2] = split
    reason: str
    confidence: float

@dataclass(frozen=True)
class Correction:
    issues: List[Issue]
    corrected: List[Interval]                # the FULL corrected set (idempotent)

class ConsistencyRule(ABC):
    name: str
    precedence: int                         # FIXED ordering; lower runs first
    @abstractmethod
    def check(self, seg: Interval, signals: "SegmentSignals") -> Optional[Issue]: ...

def reconcile(segments, signals, rules) -> Correction:    # fixed precedence; idempotent
    ...

```

5.2 Worked rules (fixed precedence)

```
| # | Rule | Evidence contradiction | Correction |
|---|---|---|---|
| 1 | `trailing_dead_time` | segment still "in rally" at its end, but the shuttle went inactive earlier | **truncate** to the last-active frame *(the owner's misfire fix)* |
| 2 | `no_net_crossing` | `net_crossings < 1` ? nothing ever crossed the net | **drop** (or flag) |
| 3 | `over_merged` | a dead-time gap **>*** `merge_gap` of shuttle-inactivity sits **inside** one segment | **split** at the gap |
| 4 | `boundary_slack` | first sustained shuttle activity / first net-crossing is well inside the start | **tighten** start to serve-contact |
```

Each rule reads only **persisted evidence** (the trajectory activity per frame, `net_crossings`, person occupancy, analyzer signals/events) ? it is a deterministic linter, not a re-detector.

5.3 Report-first, apply-optional (non-negotiable)

`reconcile()` returns **both** an issues report **and** a corrected set. It **NEVER** silently rewrites (auditability + the repo's honesty/attribution working agreement). The pipeline either:

- **(default)** persists the report and **surfaces** it for human review ? production output unchanged; or
- **(flagged, `reconcile.apply`)** applies the corrected set, **recording every change** in the corrections record (before?after, rule, evidence). Rollback = flip the flag.

It is a `Variant` transformation, so its lift is A/B-measurable BEFORE auto-apply. A reconciler is a pure post-`predict` stage: `reconcile(predict(...), signals, rules).corrected`. Wrapping that as a `RECONCILED` variant lets `compare(CONTROL, RECONCILED)` report the LOVO lift on the 6 golden videos (`QUALITY_ITERATIONS.md` corpus) ? promote to auto-apply **only** on measured lift.

Idempotent; fixed rule precedence. `reconcile(reconcile(x)) == reconcile(x)` (truncating an already-truncated segment is a no-op); precedence is a fixed table so the result is deterministic and order-independent of input segment order.

6. No-regression discipline (identical to the harness contract)

The same four guarantees `EXPERIMENT_HARNESS.md` §6 makes ? applied to both seams:

- **Default OFF.** Analyzers gate on `analyzer_flags` (default empty); the reconciler is report-first (apply gate default off). Merged code **cannot** change served behavior until explicitly enabled.
- **Shadow-persist, then promote.** Signals/issues are persisted even while ignored, so a `Variant` re-scores them offline; a signal becomes a *decision input* only when a Variant that references it wins LOVO. (This is `EXPERIMENT_HARNESS` §5.2's shadow mode for detection-touching cues.)
- **Promotion only via the eval gate.** `run_regression.py` against the golden `baseline.json` (exit 1 blocks); the experiment leaderboard records the spec hash + per-video + LOVO deltas + label-set hash.
- **Rollback = flip a config flag, never `git revert`.** The pinned `CONTROL` variant is always present; turning an analyzer off / the reconciler back to report-only is one config edit.

The keystone test (extends the harness's existing invariant): ****all analyzer flags OFF + reconciler report-only ? candidate rows AND predictions byte-identical to control.****

7. How it composes with fusion P2

- ****Analyzers plug into fusion P2's enrichment hook.**** ``FusionSegmenter._enrich_candidate_stats`` (``fusion_hybrid.py:70``) is the exact injection point §2.3 generalizes; ``net_crossings`` and the person features become the first registered analyzers (no behavior change). New cues are add-a-producer, per ADR-007 / MULTI_SIGNAL_FUSION_PLAN §2.
- ****Reconcilers consume fusion P2's persisted signals.**** They read the same ``candidate_segments.stats`` (``net_crossings``, ``players_in_court``, ?) plus the persisted trajectory ? no new detection.
- Both inherit fusion P2's control-reproducing defaults, so the fusion segmenter with everything off still seeds and hands off identically to ``wasb_hybrid``.

8. Trade-offs & risks

- ****Deterministic vs learned signals.**** An analyzer may emit a verdict-shaped signal (``service_fault=1.0``), but we persist it as a ****feature****, so the scorer still arbitrates. Risk: a confidently-wrong deterministic cue. Mitigation: the ****confidence column**** + the ****LOVO gate**** ? a bad signal earns ? 0 weight and never promotes.
- ****Ordering / idempotence (reconcilers).**** Multiple rules touching one segment must commute to a fixed point. Mitigation: fixed precedence table + an idempotence test (``reconcile?reconcile == reconcile``).
- ****Signal, not verdict.**** The whole spine. An analyzer that **decides** (drops a window) re-introduces the special-casing we are removing; reviews must reject "analyzer returns a boolean keep/drop."
- ****Report-first auditability.**** Reconcilers never silently mutate output; every applied change is recorded. This is the honesty/attribution working agreement encoded in the data path.
- ****Small-N feature discipline.**** Each analyzer adds ****few**** scale/translation-invariant columns (counts/ratios/presence, never raw coords ? MULTI_SIGNAL_FUSION_PLAN §3.2). Too many columns overfit the still-small golden set; the model stays a logistic regression, not a net.
- ****Cost & licensing.**** Pure-trajectory analyzers (service-fault) are GPU-free; person-dependent analyzers pull torch and must gate behind the person cue flag. All dependencies stay MIT/Apache/BSD (ADR-002/005).
- ****Event-table timing.**** Analyzer events precede confirmation; the buffer-then-flush wiring (§4) is the one non-obvious integration detail ? call it out in review.

9. Phased execution (owner-gated; ONE PR per slice, off `master`, default OFF)

- ****A0 ? this doc****, indexed in ``docs/README.md``, cross-linked from the peers + ADR-007. **(this deliverable; docs-only.)**
- ****A1 ? analyzer contracts + registry**** (``WindowAnalyzer`/`AnalyzerResult`/`Signal`/`AnalyzerEvent`/`BoundaryHint` + `analyzer_registry`); generalize `_enrich_candidate_stats` to iterate enabled analyzers (none enabled ? byte-identical). Retrofit `net_crossings` + person features as the first two`

analyzers

behind their existing flags. Pure, unit-tested.

- **A2** ? service-fault analyzer behind `analyzer_flags.service_fault`` (default OFF); emits the signal + `net_hit`` event + end boundary hint; shadow-persisted. Offline `compare(CONTROL, +service_fault)`` LOVO on the 6 golden videos; promote a Variant only on measured lift.
- **A3** ? event flush wiring (buffer analyzer events on candidates ? `add_event`` on confirmation).
- **R1** ? reconciler contracts + registry + `reconcile()` with `(ConsistencyRule`/Issue`/Correction`)` with the four worked rules; pure, idempotent, unit-tested; **not** wired to apply (report-only).
- **R2** ? reconciler as a `Variant`` transform: optional post-`predict`` stage; `compare(CONTROL, RECONCILED)`` LOVO on the golden set; report-only in production.
- **R3** ? corrections persistence + opt-in apply: JSON sidecar first, then the `corrections`` table `(user_version < 3`)`; apply behind `reconcile.apply``; **rollback** = flip the flag. **Auto-apply** promoted only on measured lift.

Each PR is gated by the no-regression invariant (\$6) + `run_regression.py``.

10. Verification

- **Offline**, runnable in this dev container (no GPU/torch/cv2): **pure** analyzer math on synthetic trajectories + person boxes (service-fault fires on a synthetic net-halt, stays silent on a clean rally); the **no-regression invariant** (all analyzer flags OFF + reconciler report-only ? candidate rows + predictions byte-identical to control); the **reconciler idempotence** test (`reconcile?reconcile == reconcile``) and fixed-precedence determinism; `compare(CONTROL, variant)`` deterministic on synthetic candidate rows.
- **Machine-side** (owner, GPU/WSL ? full pipeline can't run here): **persist** analyzer signals on a real clip; `compare(CONTROL, +service_fault)`` and `compare(CONTROL, RECONCILED)`` LOVO on the 6 golden videos; confirm the owner's **trailing-dead-time misfire** truncates correctly and the **held-shuttle / net-hit** service feed **is** down-weighted; spot-check `net_hit`` events landed on confirmed segments.
- **Docs** (A0): links resolve; indexed in `docs/README.md``; cross-linked from `[MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md)`, `[EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)`, `[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md)`, and ADR-007.

34. assistant (2026-06-13T19:22:25.519Z)

Now the README cross-link. I'll add it right after the EXPERIMENT_HARNESS and QUALITY_ITERATIONS entries since it's a peer to those.

35. user (2026-06-13T19:22:30.699Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\README.md has been updated successfully. (file state is current in your context ? no need to Read it back)

36. assistant (2026-06-13T19:22:38.279Z)

Let me verify referenced doc links resolve and confirm I only stage my two files.

37. user (2026-06-13T19:22:39.614Z)


```
M docs/README.md
?? docs/ANALYZERS_AND_RECONCILERS.md
---line count---
385 docs/ANALYZERS_AND_RECONCILERS.md
216 docs/EXPERIMENT_HARNESS.md
601 total
```

38. assistant (2026-06-13T19:22:56.559Z)

The doc is 385 lines vs the 216-line peer ? too long for "size/shape of EXPERIMENT_HARNESS.md."
Let me tighten it, primarily by compressing the contract code blocks and trimming verbose prose,
targeting the ~216?296 band of its peers.

39. user (2026-06-13T19:23:12.519Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\ANALYZERS_AND_RECONCILERS.md
has been updated successfully. (file state is current in your context ? no need to Read it back)

40. user (2026-06-13T19:23:17.256Z)

```
373 docs/ANALYZERS_AND_RECONCILERS.md
```

41. assistant (2026-06-13T19:23:27.984Z)

Still long. Let me re-read the current file to find the fat to trim.

42. user (2026-06-13T19:23:28.630Z)

```
1      # Analyzers & reconcilers ? two extensibility seams for rally-detection quality
2
3      **Status:** **DESIGN ? owner-gated, NOT started.** A peer to
4      [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) and
5      [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md).
6      This doc specs two *seams* ? a forward one (injectable window analyzers) and a backward
7      one (a
8      reconciliation pass) ? that let us keep piling on clever rally signals **without the
9      pipeline becoming a
10     pile of special cases.** Docs-only; no implementation code. Each phase below is a
11     separate, owner-gated PR.
12
13     > **Companion docs:** the fusion seam these plug into is
14     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md)
15     > (`RallyCue` = swappable producer; the `_enrich_candidate_stats` hook; the
16     scale/translation-invariant
17     > feature discipline, §3.2); the no-regression A/B + LOVO measurement these ride is
18     > [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) (`Variant`, `compare`,
19     `compare_unlabeled`); the scorer
20     > is the distilled logistic regression (**ADR-004**, `backend/eval/classifier.py`); the
21     modular rally
22     > layer is **ADR-007**
23     ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the 2-layer
24     > mental model is [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md); the live
25     quality log is
26     > [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md).
27
28     ---
29
30     ## 1. The spine ? NO special-casing in the decision path
31
32     This is the one principle everything else serves:
33
34     > **Every analyzer and reconciler emits a SIGNAL carrying a value + a
35     confidence/presence. The SCORING
36     > MODEL consumes those signals and *learns to weight them*. We do NOT add hand-coded
```

```

`if/else` branches
26     > into the decision.**
27
28     **Analyzers propose; the scorer arbitrates; the experiment harness + golden set decide
what gets
29     promoted.** That separation is what lets us add a net-hit service fault, a let-cord, a
double-bounce, a
30     shuttle-out cue ? each as "register a producer + a feature column," never "add a
branch." A signal
31     that doesn't earn its keep stays at weight ? 0; it costs nothing because it is an unused
column, not a
32     live `if`.
33
34     The pattern already exists in miniature. `experiment.predict()`
(`backend/eval/experiment.py:203`)
35     treats `net_crossings` and `players_in_court` as **feature columns** ? `_gate_mask` is
the trivial scorer
36     (threshold a column), `LogisticRegression` is the richer learned scorer (weight all
columns), and a
37     future scorer is richer still. The decision path is *signals ? scorer*, with **zero**
`if service_fault:
38     drop`. Analyzers feed that path more columns; reconcilers run *outside* it as an audited
post-pass.
39
40     ...
41     ???????????? DETECTION (expensive, GPU/WSL, run ONCE) ????????????
42     frames ?? WASB (frozen, MIT) ?? shuttle (x,y,v,vis)/frame ??
43     frames ?? PersonDetector (ours)?? boxes + tracks/frame ?????? per-frame-aligned
44     audio ?? hit detector (ours) ?? hit events ?????????????????? signals (+ net
estimate,
45     ?????????????????????????????????????????????????????????????? window bounds)
46     ?
47     shuttle-seeded candidate windows (FusionSegmenter, P2)
48     ?
49     ???????????? ANALYZERS ? forward seam, run PER WINDOW (flag-gated, default OFF)
????????????
50     ? WindowAnalyzer registry: service-fault · double-bounce · let-cord · shuttle-out ·
? ?
51     ? each ? AnalyzerResult{ signals(value+confidence), events, optional boundary hints
} ?
52
53     ?????????????????????????????????????????????????????????????????????????????????????????
54     ? signals merged into candidate_segments.stats
55     ?
56     DECISION / SCORING ? signals ? learned scorer (NO branches)
57     _gate_mask (trivial) + LogisticRegression (ADR-004) + harness-pinned model
58     ? predicted rally segments
59     ?
60     ???????????? RECONCILERS ? backward seam, "linter for rally decisions"
????????????????????
61     ? ConsistencyRule registry (fixed precedence, idempotent): decision-vs-evidence
conflicts ?
62     ? ? Issue + corrected set. REPORT-FIRST: report always; apply only under flag /
review. ?
63
64     ?????????????????????????????????????????????????????????????????????????????????????????
65     ?
66     candidate_segments.stats · events table · corrections record · highlights
67     ?
68     EXPERIMENT HARNESS (compare / LOVO / golden set) ? gates EVERY promotion
69     ...
70
71     ---
72
73     ## 2. Seam 1 ? injectable window analyzers (forward)

```

```

72
73     A registry of pluggable `WindowAnalyzer`s ? mirroring `segmenter_registry`
74     (`backend/pipeline/segmenters/base.py:35`) ? that run over a candidate window's **per-
frame-aligned
75     signals, all already produced**: the WASB shuttle trajectory (`frame/.visible/.x/.y` +
velocity), the

```

43. assistant (2026-06-13T19:23:41.730Z)

Let me read the reconciler and trade-offs sections to find the remaining trims.

44. user (2026-06-13T19:23:43.007Z)

```

170     ---
171
172     ## 3. How signals reach the scorer ? columns, not branches
173
174     The persisted `stats` columns are exactly what the harness already lifts.
`load_video_candidates(...,
175     feature_names=[...])` (`experiment.py:429`) reads each name out of the row's `stats`
JSON into a feature
176     column; `predict()` (`experiment.py:203`) gates/scores on those columns and merges. So a
new analyzer
177     signal becomes a scorer input by **adding its column name to a `Variant`'s
`feature_names`** ? no new
178     scoring code, no branch:
179
180     ```python
181     # A learned Variant that lets the scorer weight the new service-fault signal:
182     fusion_plus = Variant(name="fusion+service_fault",
183                           feature_names=[*SHUTTLE_FEATURES, "person__players_in_court",
184                                           "service_fault__service_fault",
185                                           "service_fault__service_fault_confidence"])
186     compare(golden_videos, CONTROL, fusion_plus)          # LOVO-honest B ? A on the 6 golden
videos
187     ```
188
189     The model **learns the weights** ? e.g. down-weight a shuttle-motion window with a high-
confidence
190     `service_fault`, up-weight a window with ?2 in-court players and multiple net-crossings.
**What the
191     classifier cannot see, it cannot special-case;** what it can see, it weights from data.
The confidence
192     column lets it discount an uncertain analyzer instead of trusting a binary verdict ? the
small-N-safe way
193     to fold in a deterministic-ish cue. (Same logistic regression as ADR-004; we add
columns, never a net.)
194
195     ---
196
197     ## 4. Persistence
198
199     - **Features ? `candidate_segments.stats`** (`database.py:240`). Namespaced
`<analyzer>__<key>`
200     (+ `_confidence`). Persisted **as a shadow** even when the served decision ignores
them, so any future
201     `Variant` re-scores offline (EXPERIMENT_HARNESS §5.2). `version`/`name` recorded for
provenance.
202     - **Events ? the `events` table** (`database.py:227`, `add_event(segment_id, ts, type,
player, details)`).
203     Its FK is to `play_segments`, which exist only **after** the AI handoff confirms a
candidate. So
204     analyzer events are **buffered on the candidate** (`stats["_events"]`) and **flushed**
to `events` in
205     Phase 4, when `_candidate_id` ? `segment_id` is linked (`trajectory_hybrid.py:283`).

```

```

Candidate-level
206     discoveries survive for offline re-scoring; the `events` table stays "events on
confirmed segments."
207     - **Corrections record** (§5) ? report-first means corrections are **recorded, never
silent.** Phase one:
208     a JSON sidecar shaped like `experiment_leaderboard.json` (one row per issue: segment,
kind, evidence,
209     before?after, reason, timestamp). When auto-apply is promoted, add a `corrections`
table via the
210     `PRAGMA user_version` migration ladder (`database.py:88`, `if version < 3:`) so
applied corrections are
211     queryable and reversible.
212
213     ---
214
215     ## 5. Seam 2 ? the reconciliation pass (backward): a linter for rally decisions
216
217     Given **predicted segments + the persisted per-frame/per-window signals**, detect where
the **DECISION
218     contradicts the EVIDENCE** and propose corrections. This is the backward dual of the
analyzers: analyzers
219     add evidence *before* the decision; reconcilers audit the decision *against* evidence
*after* it.
220
221     ### 5.1 Contracts (proposed `backend/pipeline/reconcilers/`)
222
223     ```python
224     @dataclass(frozen=True)
225     class Issue:
226         segment_index: int
227         kind: str                                # rule id, e.g. "trailing_dead_time"
228         evidence: Dict[str, Any]                 # the signal values that triggered it (auditable)
229         before: Interval
230         after: Optional[List[Interval]]          # None = drop · [iv] = corrected · [iv1, iv2] =
split
231         reason: str
232         confidence: float
233
234     @dataclass(frozen=True)
235     class Correction:
236         issues: List[Issue]
237         corrected: List[Interval]                # the FULL corrected set (idempotent)
238
239     class ConsistencyRule(ABC):
240         name: str
241         precedence: int                          # FIXED ordering; lower runs first
242         @abstractmethod
243         def check(self, seg: Interval, signals: "SegmentSignals") -> Optional[Issue]: ...
244
245     def reconcile(segments, signals, rules) -> Correction:    # fixed precedence; idempotent
246     ...
247     ```
248
249     ### 5.2 Worked rules (fixed precedence)
250
251     | # | Rule | Evidence contradiction | Correction |
252     |---|---|---|---|
253     | 1 | `trailing_dead_time` | segment still "in rally" at its end, but the shuttle went
inactive earlier | **truncate** to the last-active frame *(the owner's misfire fix)* |
254     | 2 | `no_net_crossing` | `net_crossings < 1` ? nothing ever crossed the net | **drop**
(or flag) |
255     | 3 | `over_merged` | a dead-time gap **>*** `merge_gap` of shuttle-inactivity sits
**inside** one segment | **split** at the gap |
256     | 4 | `boundary_slack` | first sustained shuttle activity / first net-crossing is well
inside the start | **tighten** start to serve-contact |

```

```

257
258     Each rule reads only **persisted evidence** (the trajectory activity per frame,
`net_crossings`, person
259     occupancy, analyzer signals/events) ? it is a deterministic linter, not a re-detector.
260
261     ### 5.3 Report-first, apply-optional (non-negotiable)
262
263     `reconcile()` returns **both** an issues report **and** a corrected set. It **NEVER
silently rewrites**
264     (auditability + the repo's honesty/attribution working agreement). The pipeline either:
265
266     - **(default)** persists the report and **surfaces it for human review** ? production
output unchanged; or
267     - **(flagged, `reconcile.apply`)** applies the corrected set, **recording every change**
in the
268         corrections record (before?after, rule, evidence). Rollback = flip the flag.
269
270     **It is a `Variant` transformation, so its lift is A/B-measurable BEFORE auto-apply.** A
reconciler is a
271     pure post-`predict` stage: `reconcile(predict(...), signals, rules).corrected`. Wrapping
that as a
272     `RECONCILED` variant lets `compare(CONTROL, RECONCILED)` report the LOVO lift on the 6
golden videos
273     ([QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) corpus) ? promote to auto-apply **only**
on measured lift.
274
275     **Idempotent; fixed rule precedence.** `reconcile(reconcile(x)) == reconcile(x)`
(truncating an
276     already-truncated segment is a no-op); precedence is a fixed table so the result is
deterministic and
277     order-independent of input segment order.
278
279     ---
280
281     ## 6. No-regression discipline (identical to the harness contract)
282
283     The same four guarantees EXPERIMENT_HARNESS.md §6 makes ? applied to both seams:
284
285     - **Default OFF.** Analyzers gate on `analyzer_flags` (default empty); the reconciler is
report-first
286     (apply gate default off). Merged code **cannot** change served behavior until
explicitly enabled.
287     - **Shadow-persist, then promote.** Signals/issues are persisted even while ignored, so
a `Variant`
288         re-scores them offline; a signal becomes a *decision input* only when a Variant that
references it wins
289         LOVO. (This is EXPERIMENT_HARNESS §5.2's shadow mode for detection-touching cues.)

```

45. assistant (2026-06-13T19:23:46.614Z)

The content is dense and all required. Let me compress the §5.1 contracts block (same treatment as §2.1) and tighten a few prose passages.

46. user (2026-06-13T19:23:53.741Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\ANALYZERS_AND_RECONCILERS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

47. assistant (2026-06-13T19:24:00.391Z)

Now tighten the §1 "pattern in miniature" paragraph and trim a couple of verbose prose lines.

48. user (2026-06-13T19:24:05.045Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\ANALYZERS_AND_RECONCILERS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

49. user (2026-06-13T19:24:08.434Z)

```
367 docs/ANALYZERS_AND_RECONCILERS.md
296 docs/MULTI_SIGNAL_FUSION_PLAN.md
216 docs/EXPERIMENT_HARNESS.md
879 total
```

50. assistant (2026-06-13T19:24:18.307Z)

367 is reasonable for a two-seam doc (the peer covers one seam in 296), and the shape matches exactly. Let me do one light trim of the trade-offs bullets, then verify links and render.

51. user (2026-06-13T19:24:19.063Z)

```
300     - **Reconcilers consume fusion P2's persisted signals.** They read the same
`candidate_segments.stats`
301     (`net_crossings`, `players_in_court`, ?) plus the persisted trajectory ? no new
detection.
302     - Both inherit fusion P2's control-reproducing defaults, so the fusion segmenter with
everything off still
303     seeds and hands off identically to `wasb_hybrid`.
304
305     ---
306
307     ## 8. Trade-offs & risks
308
309     - **Deterministic vs learned signals.** An analyzer may emit a verdict-shaped signal
(`service_fault=1.0`),
310     but we persist it as a **feature**, so the scorer still arbitrates. Risk: a
confidently-wrong
311     deterministic cue. Mitigation: the **confidence column** + the **LOVO gate** ? a bad
signal earns ? 0
312     weight and never promotes.
313     - **Ordering / idempotence (reconcilers).** Multiple rules touching one segment must
commute to a fixed
314     point. Mitigation: fixed precedence table + an idempotence test (`reconcile?reconcile
== reconcile`).
315     - **Signal, not verdict.** The whole spine. An analyzer that *decides* (drops a window)
re-introduces the
316     special-casing we are removing; reviews must reject "analyzer returns a boolean
keep/drop."
317     - **Report-first auditability.** Reconcilers never silently mutate output; every applied
change is
318     recorded. This is the honesty/attribution working agreement encoded in the data path.
319     - **Small-N feature discipline.** Each analyzer adds **few** scale/translation-invariant
columns
320     (counts/ratios/presence, never raw coords ? MULTI_SIGNAL_FUSION_PLAN §3.2). Too many
columns overfit the
321     still-small golden set; the model stays a logistic regression, not a net.
322     - **Cost & licensing.** Pure-trajectory analyzers (service-fault) are GPU-free; person-
dependent analyzers
323     pull torch and must gate behind the person cue flag. All dependencies stay
MIT/Apache/BSD (ADR-002/005).
324     - **Event-table timing.** Analyzer events precede confirmation; the buffer-then-flush
wiring (§4) is the
325     one non-obvious integration detail ? call it out in review.
326
327     ---
328
329     ## 9. Phased execution (owner-gated; ONE PR per slice, off `master`, default OFF)
330
```

```

331 - **A0 ? this doc**, indexed in `docs/README.md`, cross-linked from the peers + ADR-007.
*(this
332     deliverable; docs-only.)*
333 - **A1 ? analyzer contracts + registry**
(`WindowAnalyzer`/`AnalyzerResult`/`Signal`/`AnalyzerEvent`/
334     `BoundaryHint` + `analyzer_registry`); generalize `_enrich_candidate_stats` to iterate
enabled analyzers
335     (none enabled ? byte-identical). Retrofit `net_crossings` + person features as the
first two analyzers
336     behind their existing flags. Pure, unit-tested.
337 - **A2 ? service-fault analyzer** behind `analyzer_flags.service_fault` (default OFF);
emits the signal +
338     `net_hit` event + end boundary hint; shadow-persisted. Offline `compare(CONTROL,
+service_fault)` LOVO on
339     the 6 golden videos; promote a Variant only on measured lift.
340 - **A3 ? event flush wiring** (buffer analyzer events on candidates ? `add_event` on
confirmation).
341 - **R1 ? reconciler contracts + registry + `reconcile()`**
(`ConsistencyRule`/`Issue`/`Correction`) with
342     the four worked rules; pure, idempotent, unit-tested; **not** wired to apply (report-
only).
343 - **R2 ? reconciler as a `Variant` transform**: optional post-`predict` stage;
`compare(CONTROL,
344     RECONCILED)` LOVO on the golden set; report-only in production.
345 - **R3 ? corrections persistence + opt-in apply**: JSON sidecar first, then the
`corrections` table
346     (`user_version < 3`); apply behind `reconcile.apply`; **rollback = flip the flag.**
Auto-apply promoted
347     only on measured lift.
348
349     Each PR is gated by the no-regression invariant ($6) + `run_regression.py`.
350
351 ---
352
353 ## 10. Verification
354
355 - **Offline, runnable in this dev container (no GPU/torch/cv2):** pure analyzer math on
synthetic
356     trajectories + person boxes (service-fault fires on a synthetic net-halt, stays silent
on a clean
357     rally); the **no-regression invariant** (all analyzer flags OFF + reconciler report-
only ? candidate
358     rows + predictions byte-identical to control); the **reconciler idempotence** test
359     (`reconcile?reconcile == reconcile`) and fixed-precedence determinism;
`compare(CONTROL, variant)`
360     deterministic on synthetic candidate rows.
361 - **Machine-side (owner, GPU/WSL ? full pipeline can't run here):** persist analyzer
signals on a real
362     clip; `compare(CONTROL, +service_fault)` and `compare(CONTROL, RECONCILED)` LOVO on
the 6 golden videos;
363     confirm the owner's **trailing-dead-time misfire** truncates correctly and the **held-
shuttle / net-hit
364     service feed** is down-weighted; spot-check `net_hit` events landed on confirmed
segments.
365 - **Docs (A0):** links resolve; indexed in `docs/README.md`; cross-linked from
366     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md),
[EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md),

```

52. assistant (2026-06-13T19:24:58.661Z)

Before committing, I need to understand the working-tree state. The initial `git status` showed `docs/README.md` was **already modified** and several fusion-P2 files are uncommitted ? that's the in-flight fusion-P2 work the task warned me about. I must not let my docs-only PR drag in that in-flight work. Let me inspect precisely what's on master vs the dirty working tree.

53. user (2026-06-13T19:25:01.514Z)

```
=== README diff vs HEAD (pre-existing + my edit) ===
diff --git a/docs/README.md b/docs/README.md
index 57c1600..d3c9835 100644
--- a/docs/README.md
+++ b/docs/README.md
@@ -17,7 +17,9 @@ Live documentation for the badminton-highlight-indexer.
- [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) ? Design (owner-gated) for fusing
the shuttle cue with a person (and later audio/pose) cue in one rally-fusion layer; picks up
ADR-007's parked person cue and the named-but-unbuilt "Gate B."
- [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Latest (v8) vendor labeling
guidelines mirrored from the collector's `docs/annotation/`, with the fusion-driven court-
calibration / singles-doubles additions marked `? ADDED (fusion)`.
- [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
-- [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (owner-gated) for A/B + shadow
experimentation on rally-detection variants with zero production regression ? experiments
are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback =
flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`,
`metrics`).
+- [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (P1 SHIPPED) for A/B + shadow + SxS
experimentation on rally-detection variants with zero production regression ? experiments
are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback =
flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`,
`metrics`); implemented in `backend/eval/experiment.py`.
+- [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) ? Living log of rally-detection quality
iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-
chronological record; terminal-state snapshots get archived under `archives/quality-
iterations/`.
+- [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated, NOT
started) for two extensibility seams: injectable window analyzers (forward ? pluggable per-
window cues that emit confidence-scored signals + events + boundary hints, e.g. a net-hit
service-fault detector) and a reconciliation pass (backward ? a report-first "linter for
rally decisions" that flags where the decision contradicts the evidence and proposes
corrections). The spine: analyzers/reconcilers emit signals, the scoring model learns their
weights ? no special-casing in the decision path. A peer to `EXPERIMENT_HARNESS.md` and
`MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).

## Practical / ongoing
- [TRACKNET_WSL_SETUP.md](TRACKNET_WSL_SETUP.md)

=== Is QUALITY_ITERATIONS on master? ===
error: pathspec 'docs/QUALITY_ITERATIONS.md' did not match any file(s) known to git
=== Is MULTI_SIGNAL_FUSION_PLAN on master? ===
docs/MULTI_SIGNAL_FUSION_PLAN.md
=== full status ===
M backend/config/models.py
M backend/pipeline/detectors/person_detector.py
M backend/pipeline/detectors/rally_gate.py
M backend/pipeline/segmenters/__init__.py
M backend/pipeline/segmenters/trajectory_hybrid.py
M docs/README.md
M tests/test_config.py
M tests/test_trajectory_hybrid_equivalence.py
?? artifacts/
?? backend/pipeline/detectors/fusion_features.py
?? backend/pipeline/segmenters/fusion_hybrid.py
?? docs/ANALYZERS_AND_RECONCILERS.md
?? docs/QUALITY_ITERATIONS.md
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
```



```
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py
?? tests/test_fusion_features.py
?? tests/test_fusion_hybrid.py
```

54. user (2026-06-13T19:29:13.258Z)

```
<tool_use_error>String to replace not found in file.
String: archives/quality-iterations/`).
- [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated, NOT
started) for two extensibility seams: **injectable window analyzers** (forward ? pluggable per-
window cues that emit confidence-scored signals + events + boundary hints, e.g. a net-hit
service-fault detector) and a **reconciliation pass** (backward ? a report-first "linter for
rally decisions" that flags where the decision contradicts the evidence and proposes
corrections). The spine: analyzers/reconcilers emit *signals*, the scoring model *learns* their
weights ? **no special-casing in the decision path**. A peer to `EXPERIMENT_HARNESS.md` and
`MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).
(note: Edit also tried swapping \uXXXX escapes and their characters; neither form matched, so
the mismatch is likely elsewhere in old_string. Re-read the file and copy the exact surrounding
text.)</tool_use_error>
```

55. assistant (2026-06-13T19:29:10.011Z)

This is the critical realization the task warned about: the entire ****in-flight fusion-P2 PR is sitting uncommitted in the working tree**** ? `models.py` (FusionConfig), `fusion\_hybrid.py`, `fusion\_features.py`, `person\_detector.detections\_per\_frame`, `QUALITY\_ITERATIONS.md`, AND pre-existing README edits (EXPERIMENT_HARNESS "P1 SHIPPED" wording + the QUALITY_ITERATIONS entry). My docs-only PR must contain ****only**** my new doc + my one README line ? not the owner's in-flight work, and I must not clobber that WIP.

My plan: remove my line from the working-tree README, back up the owner's pure-WIP README, reset README to master, add my single clean line, commit my two files, then restore the owner's WIP exactly. Let me start by removing my line so the working tree returns to pure owner-WIP.

56. assistant (2026-06-13T19:29:17.235Z)

The exact text didn't match (likely the em-dashes/quote glyphs). Let me read the precise lines.

57. user (2026-06-13T19:29:17.820Z)

```
19      - [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
20      - [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (P1 SHIPPED) for A/B + shadow
+ SxS experimentation on rally-detection variants with **zero production regression** ?
experiments are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift,
rollback = flip a flag). Composes the existing `backend/eval/` machinery (`calibration`,
`regression`, `metrics`); implemented in `backend/eval/experiment.py`.
21      - [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) ? **Living log** of rally-detection
quality iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-
chronological record; terminal-state snapshots get archived under `archives/quality-
iterations/`.
22      - [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated,
NOT started) for two extensibility seams: **injectable window analyzers** (forward ? pluggable
per-window cues that emit confidence-scored signals + events + boundary hints, e.g. a net-hit
service-fault detector) and a **reconciliation pass** (backward ? a report-first "linter for
rally decisions" that flags where the decision contradicts the evidence and proposes
corrections). The spine: analyzers/reconcilers emit *signals*, the scoring model *learns* their
```

weights ? **no special-casing in the decision path**. A peer to `EXPERIMENT\_HARNESS.md` and `MULTI\_SIGNAL\_FUSION\_PLAN.md` (ADR-007 modular rally layer).

23

24 ## Practical / ongoing

58. assistant (2026-06-13T19:31:17.650Z)

My line is line 22, a clean addition after the in-flight lines 20-21. To keep my PR free of the owner's WIP, I'll isolate my README change with git plumbing (build `master + my line`, stage that, leave the owner's working-tree WIP byte-intact). First, diagnostics: confirm encoding/EOL so my reconstruction is byte-exact, and locate the anchor line in master.

59. user (2026-06-13T19:31:24.193Z)

```
working-tree README: BOM=False CRLF=True bytes=4552
master README (git): bytes=3472 (no trailing-newline info)
master EXPERIMENT_HARNESS at index 19 (line 20): -
[EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (o
master line count: 30
my ANALYZERS line occurrences in working tree: 1
```

60. assistant (2026-06-13T19:32:13.448Z)

Key finding: the working-tree README is **CRLF** ? that's why my earlier multi-line Edit (using `\\n`) failed, not the em-dashes. I'll do the README isolation with CRLF-aware PowerShell plumbing so the owner's WIP is restored byte-exact. First, capture my exact line, build a byte-exact owner-WIP-only backup, and check git's autocrlf so the committed blob matches repo convention ? without modifying README yet.

61. user (2026-06-13T19:32:25.194Z)

```
EOL=13,10 origBytes=4552 ownerBytes=3841 myLineBytes=709 delta=711 (expect 711)
autocrlf = true
master line count 30
backups written: wip=C:\Users\avidu\AppData\Local\Temp\bhi_readme_wip.bak
owner=C:\Users\avidu\AppData\Local\Temp\bhi_readme_owner.bak
myline=C:\Users\avidu\AppData\Local\Temp\bhi_my_readme_line.txt
```

62. assistant (2026-06-13T19:32:44.786Z)

Byte-exact backups verified (delta 711 = 709 + CRLF). `autocrlf=true`, so writing CRLF and `git add` yields an LF blob matching repo convention. Now build `master + my line`, stage only my two files, and verify the staged diff is exactly my one line before committing.

63. user (2026-06-13T19:32:53.906Z)

warning: in the working copy of 'docs/ANALYZERS_AND_RECONCILERS.md', LF will be replaced by CRLF the next time Git touches it

=== staged diff for README (should be ONLY my one added line) ===

```
diff --git a/docs/README.md b/docs/README.md
```

```
index 57c1600..08e1431 100644
```

```
--- a/docs/README.md
```

```
+++ b/docs/README.md
```

```
@@ -18,6 +18,7 @@ Live documentation for the badminton-highlight-indexer.
```

```
- [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Latest (v8) vendor labeling
guidelines mirrored from the collector's `docs/annotation/`, with the fusion-driven court-
calibration / singles-doubles additions marked `? ADDED (fusion)`.
```

```
- [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
```

```
- [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (owner-gated) for A/B + shadow
experimentation on rally-detection variants with **zero production regression** ? experiments
are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback =
flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`,
`metrics`).
```

```
+ [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated, NOT
```

started) for two extensibility seams: ****injectable window analyzers**** (forward ? pluggable per-window cues that emit confidence-scored signals + events + boundary hints, e.g. a net-hit service-fault detector) and a ****reconciliation pass**** (backward ? a report-first "linter for rally decisions" that flags where the decision contradicts the evidence and proposes corrections). The spine: analyzers/reconcilers emit ***signals***, the scoring model ***learns*** their weights ? ****no special-casing in the decision path****. A peer to ``EXPERIMENT_HARNESS.md`` and ``MULTI_SIGNAL_FUSION_PLAN.md`` (ADR-007 modular rally layer).

```
## Practical / ongoing
- [TRACKNET_WSL_SETUP.md](TRACKNET_WSL_SETUP.md)
=== status: staged vs unstaged ===
M backend/config/models.py
M backend/pipeline/detectors/person_detector.py
M backend/pipeline/detectors/rally_gate.py
M backend/pipeline/segmenters/__init__.py
M backend/pipeline/segmenters/trajectory_hybrid.py
A docs/ANALYZERS_AND_RECONCILERS.md
M docs/README.md
M tests/test_config.py
M tests/test_person_detector.py
M tests/test_rally_gate.py
M tests/test_trajectory_hybrid_equivalence.py
?? artifacts/
?? backend/pipeline/detectors/fusion_features.py
?? backend/pipeline/segmenters/fusion_hybrid.py
?? docs/QUALITY_ITERATIONS.md
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py
?? tests/test_fusion_features.py
?? tests/test_fusion_hybrid.py
```

64. assistant (2026-06-13T19:33:02.824Z)

The staged README diff is ****exactly my one line**** (clean against master, no CRLF noise, no in-flight changes), and only my two files are staged (``A docs/ANALYZERS_AND_RECONCILERS.md``, ``M docs/README.md``) ? every fusion-P2 file remains unstaged. Committing now.

65. user (2026-06-13T19:33:11.095Z)

```
[docs/analyzers-and-reconcilers 0f673ff] docs: design two rally-detection extensibility seams
(analyzers + reconcilers)
2 files changed, 368 insertions(+)
create mode 100644 docs/ANALYZERS_AND_RECONCILERS.md
=== commit result ===
0f673ff docs: design two rally-detection extensibility seams (analyzers + reconcilers)
0f673ff docs: design two rally-detection extensibility seams (analyzers + reconcilers)
docs/ANALYZERS_AND_RECONCILERS.md | 367 +++++
docs/README.md | 1 +
2 files changed, 368 insertions(+)
```

66. assistant (2026-06-13T19:33:17.125Z)

Commit is clean ? exactly 2 files (367 + 1 insertions). Now restore the owner's working-tree README to its byte-exact session-start state (WIP only, no my line) so the in-flight fusion-P2 work is left untouched.

67. user (2026-06-13T19:33:24.171Z)

```
README restored byte-exact to owner WIP: True (bytes=3841)
working-tree README still contains my line (should be False): False
=== in-flight fusion files must still be present + modified/untracked ===
  M backend/config/models.py
  M docs/README.md
?? backend/pipeline/detectors/fusion_features.py
?? backend/pipeline/segmenters/fusion_hybrid.py
?? docs/QUALITY_ITERATIONS.md
=== my committed doc on branch HEAD ===
HEAD has the new doc: yes
```

68. assistant (2026-06-13T19:34:08.375Z)

README restored byte-exact to the owner's session-start WIP (3841 bytes, no my line), all in-flight fusion-P2 files intact, and HEAD carries the new doc. The `M docs/README.md` now showing is purely the owner's own in-flight WIP ? untouched by me.

The task asks for a docs-only PR (do not merge). Let me push and open it.

69. user (2026-06-13T19:34:11.578Z)

```
origin https://github.com/avidullu/badminton-highlight-indexer.git (fetch)
origin https://github.com/avidullu/badminton-highlight-indexer.git (push)
---
remote:
remote: Create a pull request for 'docs/analyzers-and-reconcilers' on GitHub by visiting:
remote: https://github.com/avidullu/badminton-highlight-indexer/pull/new/docs/analyzers-
and-reconcilers
remote:
branch 'docs/analyzers-and-reconcilers' set up to track 'origin/docs/analyzers-and-reconcilers'.
To https://github.com/avidullu/badminton-highlight-indexer.git
* [new branch] docs/analyzers-and-reconcilers -> docs/analyzers-and-reconcilers
```

70. assistant (2026-06-13T19:34:14.519Z)

Branch pushed. Now opening the PR with a body that flags the docs-only/owner-gated nature and the merge-ordering dependency on the in-flight fusion-P2 work.

71. user (2026-06-13T19:34:35.076Z)

<https://github.com/avidullu/badminton-highlight-indexer/pull/116>

72. assistant (2026-06-13T19:34:39.243Z)

PR opened: [#116](<https://github.com/avidullu/badminton-highlight-indexer/pull/116>). Let me verify it contains exactly my two files and nothing from the in-flight fusion work.

73. user (2026-06-13T19:34:42.327Z)

```
=== PR #116 files (must be exactly 2 docs) ===
{"base":"master","files":["docs/ANALYZERS_AND_RECONCILERS.md","docs/README.md"],"head":"docs/analyzers-and-reconcilers"}

=== working tree: my doc committed, owner WIP untouched ===
  M backend/config/models.py
?? backend/pipeline/detectors/fusion_features.py
```

```
?? backend/pipeline/segmenters/fusion_hybrid.py
?? docs/QUALITY_ITERATIONS.md
?? tests/test_fusion_features.py
?? tests/test_fusion_hybrid.py
```

74. assistant (2026-06-13T19:34:58.907Z)

Verified end-to-end:

```
- **PR [#116](https://github.com/avidullu/badminton-highlight-indexer/pull/116)** ? base
`master`, contains **exactly** `docs/ANALYZERS_AND_RECONCILERS.md` + the 1-line `docs/README.md`
cross-link.
- The entire in-flight **fusion-P2 WIP is byte-intact** (`models.py`, `fusion_features.py`,
`fusion_hybrid.py`, `QUALITY_ITERATIONS.md`, tests) ? I isolated my README line via git plumbing
and restored the owner's working-tree README to its exact session-start bytes.
```

This is a meaningful checkpoint (PR pushed) ? per the working agreement I'll refresh the live current-state memory. Let me read it first.

75. user (2026-06-13T19:34:59.696Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-463 of 1603 total (73529 tokens, cap 25000). Call Read with offset=464 limit=463 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-reminder>

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-13 (RALLY-QUALITY TRACK: HARNESS + PERSON-CUE EXTRACTION SHIPPED
? **PR #114 + #115 MERGED**,
13     master `df7d5df`; suite 421?441 green; 0 open PRs). **PR #115 (fusion P1):** extracted
the torchvision person detector
14     + ByteTrack + velocity windowing out of `yolo_hybrid` into
`backend/pipeline/detectors/person_detector.py`
15     (`PersonDetector`, a reusable cue producer); `yolo_hybrid` now ORCHESTRATES only
(chunk/prefetch/AI-handoff/DB) and
16     delegates Stage-1 to `self.person`. **No behavior change** ? proven by the unchanged
`process_video` output tests (mock
17     seam moved to `self.person`); pure detection unit tests ? new
`tests/test_person_detector.py`. person_detector is
18     mypy-CLEAN (`self.model: Any`) so it needed NO quarantine; the torchvision code that put
yolo_hybrid in quarantine left
19     it, but yolo_hybrid STAYS quarantined for PRE-EXISTING orchestration debt (implicit-
Optional process_video defaults,
20     Success|Failure ABC-override mismatch, a Future[...] annotation ? flagged in mypy.ini to
ratchet separately). ruff+mypy
21     clean, CI 4/4. Docs synced (fusion-plan P1?DONE, CODE_MAP person_detector entry +
yolo_hybrid-as-orchestration, NEXT_STEPS
22     item 2 ticked). ?? **NEXT (owner-steered):** **PR 3 = fusion P2.** The cheap, highest-
confidence, PURE-LOGIC piece =
23     **wire Gate A (`min_crossings`) into the live trajectory path**
(`trajectory_hybrid`/`wasb_hybrid`) + add the knob to
24     `IndexingConfig` ? *fixes the owner's held-shuttle false positive*, default OFF,
immediately A/B-able via the harness on
25     the 6 golden trajectories. Gate B (court-mask + player-count) is **owner-gated** on the
court-mask-source decision
```

26 (MULTI_SIGNAL_FUSION_PLAN §9.1: manual polygon vs auto homography ? owner). **P3 learned fusion + the SxS-on-new-videos

27 RUN need GPU/owner-side** (persist person features on real video ? train on the 6 golden ? `compare_unlabeled(CONTROL,

28 +person)` on new footage). Two PRs this session within [[working-agreement-merging]] self-merge authority.

29

30 **Checkpoint:** 2026-06-13 (EXPERIMENT HARNESS P1 SHIPPED ? **PR #114 MERGED**, master `ecee786`; suite 421?440 green).

31 Picked up the rally-quality track (the desktop-agent pick-up `NEXT\_STEPS` flagged; both design docs `EXPERIMENT\_HARNESS.md`

32 + `MULTI\_SIGNAL\_FUSION\_PLAN.md` arrived via the session-start `git pull` 15e4e0e?9231af2, PR #112). Built

33 `backend/eval/experiment.py` = the thin variant/experiment layer from `EXPERIMENT\_HARNESS.md` P1 (the audit's "-80%

34 already exists" HELD ? verified file:line; pure glue over metrics/calibration/classifier/training/regression.gt_hash/

35 gemini_refine.merge_overlaps/query_candidate_segments, **NO new scoring math**): `Variant` (JSON recipe:

36 segmenter+config_overrides+cue_flags+Gate-A `min\_crossings`+Gate-B `min\_players`+learned classifier; `spec\_hash`; pinned

37 `CONTROL`), `compare` (labeled A/B, **LOVO-honest** train-on-others for learned variants, B?A deltas + `label\_set\_hash`),

38 **`compare\_unlabeled`** (the owner's **SxS-on-NEW-videos** ask ? agreed/a_only/b_only via match_intervals, no GT),

39 `record\_experiment` (? `eval\_baselines/experiment\_leaderboard.json`), `load\_video\_candidates` (DB bridge). 19 tests incl.

40 the load-bearing **no-regression invariant** (all cues OFF ? byte-identical to CONTROL); ruff+mypy clean; CI 4/4. Docs

41 synced (EXPERIMENT_HARNESS P1?DONE, CODE_MAP entry + stale regression-baseline-path fix, NEXT_STEPS item 3 ticked). Owner

42 answered ramp-up Qs: **corpus = all 6** golden videos (he said "5"; LargeTestVideo/TestLargeVideo may be cuts of one match

43 ? kept all 6 per the manifest); **merge authority = self-merge on green then continue to PR 2** (within [[working-agreement-merging]]).

44 ?? NEXT (this session): **PR 2 = fusion P1** person/vision-cue extraction (lift torchvision+ByteTrack+velocity out of

45 `yolo\_hybrid` ? `backend/pipeline/detectors/person\_detector.py` producer; yolo_hybrid consumes it, no behavior change ?

46 regression-safe). Then P3 learned fusion trains on the 6 golden videos + `compare\_unlabeled(CONTROL,+person)` SxS on new

47 footage (GPU/owner-side). Tracks [[candidate-segments-finetuning]] + [[collector-indexer-bridge]].

48

49 **Checkpoint:** 2026-06-13 (WORKING-TREE CLEANUP ? master `15e4e0e`, tree now CLEAN). Cleaned up the leftover

50 uncommitted bits (the cache-hygiene CODE was already MERGED as #109/#110 per the checkpoint below ? nothing to rescue

51 there): (1) `config.json` GPU-run knobs **reverted to committed defaults** (was wasb_hybrid/7200/8 ? gemini/1800/1;

52 values stay recorded in the cache-hygiene checkpoint below ? re-apply `default\_segmenter: wasb\_hybrid` in one line if

53 that's the intended new default; user didn't answer keep-vs-revert, took the safe recoverable path). (2) stray root

54 `arial.ttf` (unused ? shipped watermark font is committed `backend/assets/fonts/RobotoSlab-Regular.ttf` from #103) ?

55 moved to `scratch/`. (3) genuine untracked draft `docs/STORAGE\_SHARING\_MODEL.md` (190-line storage-anchored

56 monetization model, owner co-design pending) ? branched + **PR #111 OPEN (owner-gated, docs-only ? do NOT auto-merge**;

57 supersedes COMMERCIALIZATION.md framing if ratified ? ADR then). Master tree now shows ONLY intentional carryovers

58 (untracked `scratch/` handoffs+repro scripts, `artifacts/` Gen-0 bundle) ? could be gitignored for a pristine

59 `git status` if desired (offered, not done). Reminder: `python -m tools.cleanup\_caches`

```

--apply` reclaims the 96.5 GB
60     the SessionStart hook flagged.
61
62     **Checkpoint:** 2026-06-13 (CACHE HYGIENE shipped ? **PR #109 + docs #110 MERGED**,
master `15e4e0e`; + the ?5 reel +
63     manual cleanup). **#110** = diligence docs follow-up (README $Cache Hygiene + knobs +
tools/ line; CODE_MAP detector
64     self-clean, tools/cleanup_caches.py entry, common-task rows; min_shot_count API-parity
note) ? docs-only, CI 4/4 green.
65     Both repos now have **0 open PRs**. After the GX010125 E2E run left ~163 GB in WSL
`~/clips`, built a full cache-hygiene system (all on master):
66     (1) **self-cleaning runners** ? `indexing.{wasb,tracknet}.keep_frames` (default false)
deletes the frame cache +
67     staged video after a SUCCESSFUL run; on failure keeps them for resume; (2)
**tools/cleanup_caches.py** dry-run-first
68     ops tool (`--apply`/`--keep-video`/`--older-than`/`--include-wasbcache`/`--summary`);
(3) **disk guard** `min_free_gb`;
69     (4) **SessionStart nudge hook** in `.claude/settings.local.json` (local, gitignored).
Suite 409? **421 green**, ruff+mypy
70     clean. Policy saved ? [[cache-hygiene-policy]]. ? WSL-scan footgun: the owner's login
shell mangles `for e in ` loops
71     (loop var empty) ? use du/find arg-globbing. Also this session (earlier): merged
**#107** (API min_shot_count) + **#108**
72     (watermark Windows colon fix); ran GX010125 E2E on GPU ? **GX010125_highlights.mp4** (11
rallies >7 shots) +
73     **GX010125_highlights_min5.mp4** (24 rallies ?5, watermark now works ? validates #108)
via a stitch-only reuse script
74     (no re-inference); confirmed **Gemini got 1080p** (4K only at final stitch); delivered
an execution flowchart. Manual
75     WSL cleanup freed ~68 GB (kept GX010125's 93 GB frames + proxy per owner choice). ?
LOCAL uncommitted `config.json`:
76     default_segmenter=wasb_hybrid / min_shot_count=8 / wasb.timeout_sec=7200 from the run.
See [[gotcha-long-runs-gpu-wsl]].
77
78     **Checkpoint:** 2026-06-13 (SYNC + HANDOFF REFRESH ? read-only reconciliation by a ramp-
up session; no new code from
79     me this pass). `git pull --ff-only` = already up to date at **master `6de70aa` (#108)**.
Reconciled remote vs my last
80     session (#100): seven PRs landed from OTHER slates (NOT my work ? facts from git/gh):
**#101** configurable branding
81     watermark . **#102** thread config stitching paddings into live VideoCompiler sites .
**#103** bundle Apache-2.0
82     default watermark font + fallback log . **#104** C13 field kit (flyer/playbook/answer-
sheet/signals) . **#105** i18n
83     design plan (Kannada-first) . **#107** enforce `stitching.min_shot_count` on
/api/compile . **#108** escape Windows
84     drive-letter colon in ffmpeg filtergraph (fixes #103 watermark on Windows). **OPEN:
#106** docs/video-locality-model
85     (awaiting owner). ? **UNCOMMITTED WIP IN TREE (sibling/owner ? DO NOT CLOBBER):** a WSL
**cache-hygiene** slice in
86     progress, built on my `WslCommandMixin` ? `base.py` (`_wsl_free_gb`/`_wsl_rm_rf`),
`wasb_runner.py`+`tracknet_runner.py`
87     (`keep_frames`/`min_free_gb` config + pre-run disk guard + post-success cache delete),
`models.py` (new fields),
88     `config.json` (the GPU-run knobs: default_segmenter=wasb_hybrid, min_shot_count=8,
wasb.timeout_sec=7200), untracked
89     `arial.ttf` (watermark font) + `docs/STORAGE_SHARING_MODEL.md`. Not yet PR'd by whoever
owns it. [[session-handoff]]
90     rewritten to this state.
91
92     **Checkpoint:** 2026-06-13 (FIRST FULL E2E GPU RUN on a real 4K GoPro match + 2 merged
fixes).
93     Owner asked to run the indexer E2E on the GPU over `F:\GoPro Hero Black
12\KhelSutra\Badminton 12
94     June 2026\GX010125.MP4` (12.8-min 4K60, 11.5GB), highlights for rallies with >7 shots,

```

```

config-driven.
95     - **DELIVERED:** `output\GX010125_highlights.mp4` ? **11 rallies (shot_count 8?16) of 52
detected**;
```

96 true 4K 3840x2160 h264 59.94fps, 2:23, 1.15GB; stitched from the 4K ORIGINAL (proxy only for indexing).

97 wasb_hybrid: 64.8% shuttle visibility, 18 candidate windows ? 52 Gemini-confirmed rallies. min_shot_count=8

98 correctly excluded a 7-shot rally (">7" semantics). Gemini ran clean (no 503s); ~47-min wall clock.

99 - **2 PRs merged on green CI** (owner's [[working-agreement-merging]] standing authorization):

100 **#107** `/api/compile` now enforces `stitching.min\_shot\_count` (parity w/ run_process_and_stitch);

101 **#108** sticher `\_ff\_quote` escapes the Windows drive-letter colon ? fixes the on-by-default

102 bundled-font watermark (PR #103) silently failing on Windows (`fontfile='C:/...'` broke ffmpeg

103 drawtext parsing; also fixed absolute logo `movie=` paths). Both with regression tests; suite 409 green.

104 - **The hard-won run recipe is in [[gotcha-long-runs-gpu-wsl]]:** Modern Standby killed the first 2

105 overnight attempts (keep-awake fix); cv2 PNG extraction (~38min) is the real bottleneck not the GPU

106 (ffmpeg pre-extract, count-exact for CFR ? skips it); 1080p-proxy-index / 4K-stitch; per-phase

107 observability (cache manifest for WASB, handoff dir for Gemini). Reusable scripts in `scratch/`

108 (run_gx010125.py, keep_awake.ps1, probe_*.py).

109 - ? LOCAL UNCOMMITTED `config.json`: default_segmenter=wasb_hybrid, min_shot_count=8, wasb.timeout_sec=7200

110 (set for this run; revert to gemini/1/1800 if not wanted). Watermark fix NOT applied to the delivered

111 reel (it ran before #108); a re-stitch would hit the DB segment cache + now-working watermark (~5min).

112

113 **Checkpoint:** 2026-06-12 (latest ? PAID COMPONENT REMOVED + VIDEO-LOCALITY MODEL; **PR #104 updated `9c84fea` +

114 PR #106 OPEN `85abe26`, both CI green, both AWAITING OWNER**). Owner spec: (1) remove ALL associate-pay from the

115 field kit ("something else in mind") ? flyer/playbook now "Engagement terms: discussed individually"; honesty

116 mechanics rewired ("judged on honest capture, not positive answers"); KEPT: WTP research questions (study subject)

117 + the ACADEMY's Phase-2 session fee (different party ? flagged, owner may strip too). (2) Storage/upload questions

118 answered from code: uploads ARE stored in full (`security.upload\_dir`, NO retention policy ? BACKLOG); 10GB CANNOT

119 upload (max_upload_mb=4096 rejects at init; ?90MB sequential parts; hours on Indian uplinks). (3) NEW

120 `docs/VIDEO\_LOCALITY\_MODEL.md` (#106): central insight = timestamps are the product, video is dead weight (only the

121 compiler touches original bytes, cuttable where the original lives); locality ladder L0?L4 (L2 proxy-upload reuses

122 the collector ADR-002 frame-identical contract; L1-direct = client?Gemini, zero bytes through us; L0 = owned-model

123 north star; pilot = L4 sneakernet); **frugality frame (owner): bills affordable iff GPU doesn't burn ? L2/L1 serving

124 is GPU-less by construction, Gemini \$/hr, GPU stays on the owned 2070S, per-video cost ceiling required**;

125 **academy-visit dividend: Phase-2 consented recordings ARE the multi-court target corpus, show-back = live L2 dress

126 rehearsal, demo reactions = build backlog**; 10 open questions OQ1?OQ10 (retention, MD5 identity, Gemini consent,

127 key custody, packaging, review source, reel delivery, uplink data, volume, L3 survival); rec lean = L2 default

128 (NOT a greenlight). NB: #106's base commit was pushed by owner/sibling mid-session
(shared checkout). ?? owner:
129 review/merge #104 + #106; decide academy-fee keep/strip; the "something else" comp
structure is his to define.
130
131 ****Checkpoint:**** 2026-06-12 (latest ? C13 KIT REFRAMED PRODUCT-FEEDBACK-FIRST per owner
spec + ****RECRUITING FLYER**
132 added; PR #104 UPDATED `985a329`, CI 4/4 green, AWAITING OWNER REVIEW** ? do NOT merge
without him). Owner spec:
133 hire someone to visit academies, talk to coaches AND students (how does it help +
improvement feedback); stretch =
134 record ? pipeline ? show the academy ITS OWN sample result; UI/hosting "ready by then".
Kit now = 4 docs by audience:
135 ****`PRODUCT\_FIELD\_ASSOCIATE\_FLYER.md`**** (NEW ? candidate-facing recruiting flyer, owner
fills ?/phone + recruits) .
136 ****`FIELD\_VISIT\_PLAYBOOK.md`**** (RENAMED from FIELD_RESEARCH_ASSOCIATE_FLYER ? day-one
playbook: demo-after-Q6 rule,
137 adult-18+-students-only w/ coach-points-out gating, Phase-2 show-back + scripted "can we
keep using it?" no-promise
138 reply) . answer sheet (+demo-reaction/student blocks) . PMF_FIELD_SIGNALS (verdict
machinery UNCHANGED; demo/student
139 data = product feedback only). 2nd adversarial critique (`w5aexakz2`, 30 agents): 12
confirmed+fixed (key honesty
140 fixes: "web app IS opening"? "we're building?will be able to try"; "match report"? "simple
summary"; "finds EVERY
141 rally" dropped; minors loophole in "senior/adult students" closed), 15 refuted. ??
owner: review #104 (thresholds
142 \$3 + flyer copy) ? merge ? fill ?/phone ? trim 60s demo reel ? recruit.
143 Owner ratified in-chat: ****PMF validation = physically visiting Bangalore academies and
reading signals.**** Built the
144 decision layer on top of the existing `docs/FIELD\_RESEARCH\_ASSOCIATE\_FLYER.md` (which IS
the "recruit hand-out" the
145 owner remembered; the khelacharya PDFs in OneDrive are the separate internship flyer ?
see
146 **[[sibling-internship-repo]]**): ****`docs/PMF\_FIELD\_SIGNALS.md`**** (H1-H7 ? script-question
map, G/A/R anchors, H5
147 counting rule = pre-ladder-or-budget-converted ??300 only ? anchored ?300 NEVER counts,
pinned stopping rule n=15
148 target/logistics-only/verdict-once, numeric verdicts PROCEED / PIVOT A record+index /
PIVOT B partner-API
149 (PB Vision*PodPlay footnote now in-repo) / PIVOT C parent-payer / FALSIFIED) +
****docs/FIELD_VISIT_ANSWER_SHEET.md`****
150 (printable per-visit form: verbatim answers, required most-negative quote, before/after-
ladder tick, OFFICE-USE
151 score grid) + flyer Q5 rewritten open-ended-first (old wording read the ?100/300/500
ladder ALOUD ? center-bias
152 manufactures the exact ?300 PROCEED threshold; caught by the 38-agent adversarial
critique `ww0rvaqf3`: 1 high/2 med/
153 2 low confirmed+fixed, 9 refuted) + Stupa added to listen-for + NEXT_STEPS item 5 = kit
COMPLETE. ?? owner: review
154 #104 thresholds ? merge ? fill flyer ?/phone placeholders ? recruit the associate ?
first 2 visits paired. PDFs of
155 flyer+answer-sheet: generate AFTER placeholders are filled (offer stands).
156
157 ****Checkpoint:**** 2026-06-12 (latest ? ROORA VENDOR LOOP COMPLETE: ****collector #23+#24+#25
MERGED**** sequentially on
158 green CI, owner pre-authorized; master `ca61405`, 0 open PRs, only owner's `.gitignore`
edit uncommitted).
159 ****#23 boundary_grid_snap BLOCKER**** in detect_issues (gcd-of-boundary-frame-diffs > ±0.3s
bar ? gate BLOCKED; skips
160 CSV-path/no-fps/<3 rallies) ? ? ****replayed against the REAL 6-June PB+TT pilot zips:
both auto-block with "All 12
161 rally boundaries sit on a 30-frame grid (~0.50s @ 59.94fps)"**** (the manual forensics,
now automated) + ROORA_BRIEF
162 Phase-B revision (\$7 REQUIRED method: CVAT step 1, first/last box on exact contact/dead

frame, sparse interior OK;

163 \$4 grid-snap auto-detected warning; \$2 skip-rallies-in-progress-at-file-edges; \$1 annotation-optimized-copy note).

164 ****#24 prep-annotation-batch****: mirrors `//` ? proxies-only tree w/ same layout

165 (thin orchestrator over the hardened single flow; per-video isolation; resumable "have" skips; live 2-court E2E

166 smoke green). ****#25 R00RA_RUNBOOK.md**** (docs/annotation/): operator E2E ? send checklist (layout?batch?manifest

167 commit?intake+email w/ the 2 non-negotiables), receive (gate decision table incl. grid-snap, review sheet,

168 import/export on ORIGINALS), troubleshooting table for every guard. Suite ****180 green****. ?? Vendor loop is now fully

169 tooled+documented. ?? AGREED NEXT (owner, session wrap 2026-06-12): (1) owner fills the canonical doc's TODO(avidullu)

170 markers + screenshots, labels it ****v8****; (2) ****joint TEST RUN** of the full vendor loop on real footage BEFORE

171 anything goes to Roora** (runbook \$1 send-side end-to-end: batch layout ? prep-annotation-batch ? manifest ?

172 mock-receive via validate-delivery on an old delivery); (3) only then send. ****v7 ANOMALY RESOLVED (2026-06-12)****

173 the canonical vendor-facing Brief+Guide is the Google Doc at docs.google.com/document/d/1NvAXE0q7t4Z0QddLkEfnta0DU982yDZaqkGNPMe8EY0

174 (provenance: generated from master `ca61405`; supersedes Drive v3/v5/v6/v7; owner confirmed superset-of-v7) ? see

175 `[[roora-guidelines-drive-doc]]`; before sending: pick ONE version label + fill its TODO(avidullu) markers/screenshots.

176

177 ****Checkpoint**** 2026-06-12 (earlier ? ADR-002 IMPLEMENTED: ****collector #22 MERGED**** `a068fd1`, owner authorized

178 merge-on-green; #21 ADR + #20 + indexer #102/#103 all merged earlier today, slate was clean). Shipped: `curate

179 prep-annotation-proxy` + `src/core/proxy.py` ? ?1080p/-g 15`/CRF 19/audio-kept re-encode w/ `-fps\_mode passthrough`

180 (NEVER `-r`), post-encode parity verify (exact frame count; fps $\pm 0.1\%$; duration like-with-like stream/container),

181 delete-on-ANY-failure (incl. Ctrl-C), VFR refusal w/o `--allow-vfr`, provenance ?

182 `data/annotation\_proxy\_manifest.csv` (source+proxy MD5s; video_id collision guard BEFORE encode; Excel-lock prints

183 row + keeps proxy), same-file guard normcase/realpath/samefile (NTFS case-variant `--out`+`--force` would have let

184 `ffmpeg -y` TRUNCATE the source ? review verifier reproduced it live pre-fix). Suite ****169 green**** (+32), CI green.

185 Verified live on ffmpeg 8.1.1 (smoke clip 59.94fps: parity 179 frames both files; smoke caught `-vsync` deprecation ?

186 `-fps\_mode`). Process: ultracode adversarial review workflow (`wyv8524zf`, 18 agents) confirmed 1 high/2 med/7 low ?

187 all fixed pre-merge. Docs: INGESTION_PIPELINE step 0 + diagram leg + onboarding FAQ + sampling-FAQ split

188 (highlights-vs-golden), ADR-002 provenance bullet as-implemented amendment, PIPELINE_ROLE ADR refs qualified.

189 ?? REMAINING ADR-002 follow-ups (separate items, NOT started): Phase-B brief revision (frame step 1 / exact boundary

190 frames / "we supply proxies") ? owner words vendor comms; grid-snap QA check in validate-delivery. First real use:

191 run prep-annotation-proxy on the next Roora batch before handoff.

192

193 ****Checkpoint**** 2026-06-12 (earlier ? ANNOTATION-PROXY DECISION RATIFIED ? ****collector PR #21 OPEN****, docs-only

194 ****ADR-002**** in collector `docs/DECISIONS.md`). Owner agreed in-chat w/ my rec from the assessment checkpoint below:

195 proxies generated in-repo via a future `curate prep-annotation-proxy`. ADR records the proxy contract (SAME fps ?

196 never reduce; no trim; frame-count parity check; 1080p/CRF18-20/-g15/audio kept; intake-manifest row w/ original+proxy

197 MD5s), the why-collector rationale (owns the vendor flow = command #0 of convert-
 cvat?validate-delivery?import-local;
 198 verification+provenance is the durable 80%, one-off scripts skip it), and the **Phase-B
 necessity argument the owner
 199 wanted captured: ~12-15 videos/sport × ~30 min 4K/59.94 = >100k frames/file ? vendor
 frame-step-1 annotation
 200 infeasible without proxies ? near-precondition, not optimization**. Follow-ups recorded
 IN the ADR: Phase-B brief
 201 convention line (step=1, first/last box on exact contact/dead frame) + re-add grid-snap
 QA to validate-delivery.
 202 Collector checkout back on master (owner's stray `.gitignore` edit untouched). ? sibling
 PR #20 (below) edits
 203 rally_cvat.py docstring rationale 0.5/1.0?0.5/0.5 ? no file overlap w/ #21, both docs-
 true. ?? owner reviews/merges
 204 #21 (+ #102/#20 below); subcommand IMPLEMENTATION stays owner-gated (separate PR on
 greenlight).
 205
 206 **Checkpoint:** 2026-06-12 (stitching-padding wiring fix ? **indexer #102 + collector
 #20 OPEN, CI green, awaiting
 207 owner review**). Audit finding (d) (2026-06-10) fixed: config.json
 `stitching.{begin,end}\_padding` was silently
 208 ignored by the live compiler ? both `backend/main.py` VideoCompiler sites
 (`/api/compile` + CLI trim) passed only
 209 `watermark=`, so reels always used ctor defaults 0.5/**1.0** vs configured 0.5/0.5. Fix
 threads the typed
 210 StitchingConfig paddings into both sites (+ `StitchingConfig` re-export from
 backend.config); regression test drives
 211 the REAL compiler through /api/compile (mocked ffmpeg) asserting a config override
 (2.0/3.0) reaches `-ss/-to`
 212 (8.0/17.0 for a 10?14s segment ? unwired defaults would give 9.5/15.0); CODE_MAP truth-
 synced (ctor-vs-config
 213 distinction + recipe row now cites the main.py wiring). ? **Behavior change on merge:
 live API reels end padding
 214 1.0?0.5s** (the config value). Companion **collector #20** = docstring-only:
 `rally\_cvat.py::detect\_issues` cited the
 215 hardcoded 0.5/1.0 + a stale module path; now cites config-driven 0.5/0.5 (rationale
 still holds; ? cross-checks the
 216 sibling Roora-proxy thread below ? boundary-slack rationale now assumes 0.5s end
 padding, not 1.0s). Verified: indexer
 217 suite **401 green** + 4/4 CI lanes; collector **137 green** + CI. Both repos back on
 master; owner's uncommitted
 218 collector `.gitignore` edit untouched. NOTE: untracked `arial.ttf` at indexer root (not
 mine ? left alone, likely a
 219 watermark-font experiment). ?? owner reviews/merges #102 + collector #20.
 220
 221 **Checkpoint:** 2026-06-12 (Roora ANNOTATION-PROXY assessment delivered ? AWAITING OWNER
 DECISION, no code).
 222 Owner asked: re-encode videos to a Roora-friendly profile before handoff? Researched via
 workflow `wr6mnel6c`
 223 (3 sweeps + adversarial verify; full JSON
 `%LOCALAPPDATA%\Temp\claude\?\tasks\wr6mnel6c.output`). **Facts:** root
 224 cause of the ?0.5s boundaries = 30-frame CVAT sampling grid (NOT playback delay;
 direction never measured ? docs say
 225 "off", not "late"); same grid was flagged [HIGH] in the BADMINTON round too
 (`data/roora\_badminton\_issues.txt`) before
 226 PB/TT; operative vendor contract = CVAT XML, times recomputed `frame/fps`
 (`rally\_cvat.py:266`), collector joins by
 227 human video_id NOT MD5 ? proxy handoff is safe if timeline-preserving. **My rec:** add
 `curate prep-annotation-proxy`
 228 to the COLLECTOR (not a loose script): 4K?1080p, SAME fps (never drop ? their 30-frame
 habit at 30fps = 1.0s grid),
 229 keep audio, dense keyframes (-g 15), frame-count parity check, manifest row w/
 original_md5+proxy_md5. Proxy alone ?
 230 fix ? must pair w/ the review's "first/last box on exact frame, step=1" convention in
 the Phase-B brief. Also: re-add

```

231     a grid-snap QA check to validate-delivery (detect_issues deliberately dropped
coarse_sampling, rationale =
232     highlight-padding absorption ? conflicts w/ golden  $\pm 0.3s$  need). Spawned chip
`task_d8df38bc` (stitcher padding config
233     not wired into VideoCompiler ? audit finding (d) re-verified live at main.py:508/633).
?? owner decides: collector PR
234     vs one-off; then draft PR on greenlight.
235
236     **Checkpoint:** 2026-06-11 (UI-alpha **PR-2 #99 MERGED + docs-sync #100 MERGED** ? owner
authorized merge w/ due
237     diligence; master `020ff61`). **LIVE E2E AT MERGE (server up on :8123, real
artifacts): ** 86MB smoke video uploaded in
238     3 sequential parts ? reassembled to **MD5 byte-identity** with the source; negatives
verified (out-of-order 409,
239     bad-ext/traversal 400, unknown-id 404; download traversal blocked at BOTH layers ?
handler 400 on backslash form,
240     router 404 on %2F form); **compile of the real confirmed rally ? 3.9MB reel streamed via
GET /api/highlights** (the
241     fixed segment_ids path works browser-equivalent). #100 = doc due diligence: NEXT_STEPS
(PR-2 merged + sweep recorded),
242     UI plan PR-2 checked off w/ E2E evidence, **BACKLOG shlex item closed** (residual
tenancy note kept). The CUJ runbook
243     (PR2_HANDOFF $"Tomorrow-morning") is now fully unblocked for the owner. ?? NEXT: owner
CUJ; PR-3 identity (remember:
244     `allowed_video_root` must contain `upload_dir`). Original handoff context below:
245     (superseded) **Checkpoint:** 2026-06-11 (UI-alpha PR-2 executed ? #99 was awaiting
review ? per the Cowork handoff
246     `scratch/PR2_HANDOFF.md`; protocol = Cowork designs, Claude Code runs mechanics, **Avi
reviews/merges ? do NOT
247     auto-merge #99**). #88 confirmed CLOSED ? PR88_FIX skipped. Cowork's uncommitted B2+B3
implementation verified to
248     layer cleanly on post-sweep master (AI-7 CODE_MAP entries + pydantic fixes intact) ?
Step 0 run here (Cowork sandbox
249     down): suite **388 green** (15 new `tests/test_upload_download.py`), mypy clean; 3
trivial ruff B904s in the new
250     endpoints fixed per the handoff's trivial-fix rule (exception chaining only, no design
changes). Branch
251     `feat/upload-download`, exactly the 8 handoff files committed ? **including the two
long-preserved doc edits**
252     (DEFERRED_HARDENING_PLAN #16-banner + UI_ALPHA_EXECUTION_PLAN working-agreement
rewrite), which are now COMMITTED
253     (tree finally clean of them). CI 4/4 green on #99. ?? NEXT: owner reviews/merges #99 ?
then the tomorrow-morning CUJ
254     runbook in the handoff (record ? localhost:8000 ? path-paste ? index ? compile ? ?
download; gemini default,
255     rotated key in .env, ~$0.05/run). PR-3 = per-user dirs + allowed_video_root?upload_dir
(hosted-mode note).
256
257     **Checkpoint:** 2026-06-11 (8?9 QUALITY SWEEP ? COMPLETE ? all approved AIs landed;
master `31069cb`). Final stretch:
258     **AI-7 hybrid-twin collapse** = shared `segmenters/trajectory_hybrid.py` engine
(tracknet/wasb now ~40-line subclasses;
259     one `_shot_count_from`; `_wsl_bash` ? `detectors/base.py::WslCommandMixin`;
DetectorRunner ABC now declares the
260     parse/windowing contract; mypy quarantine ratcheted ?2; 13 equivalence tests; CODE_MAP
2.1/2.2 updated). ? PROCESS SLIP
261     (disclosed): AI-7 commit `2c241a8` went DIRECTLY to master (post-#96 script left
checkout on master) ? content was
262     pre-approved + locally verified; master CI green post-hoc; chose NOT to force-rewind
(concurrent sessions live).
263     - **QUALITY CHECKPOINT (owner-mandated) PASSED:** offline LOVO reproduced **EXACTLY
0.626** (Mahadevpura Gen-0 0.744)
264     post-sweep ? featurizer/eval edits numerically inert. **Live smoke = wasb_hybrid's
FIRST-EVER verified E2E run**
265     (3-min Mahadevpura slice): P1 healthcheck?P2 trajectory (9 candidates, correct

```

```

video_id keying)?P3 Gemini (503
266 capacity wave; 6-attempt backoff worked; 1 segment confirmed)?P4 add_play_segment +
candidate link
267 (`wasb-hybrid-gemini-flash-latest`). 1/9 confirms = Gemini weather, NOT code.
268 - ? Smoke found 2 REAL pre-existing live-path bugs (masked: trajectories were always
hand-scripted, tests mock WSL) ?
269 **#97 MERGED**: (1) shlex.quote strangled `~` expansion in staging (the BACKLOG
"shlex-quote breaks ~" item ? CLOSED;
270 fix = `wsl_tilde_quote`, quotes only after `~/`); (2) relative `output/wasb` resolved
against WSL cwd ? 8-min GPU run
271 wrote CSV inside WSL invisibly (fix = abspath before /mnt translation). Regression
tests pin both command shapes.
272 - Merged this sweep: **#90** hygiene . **#91** pydantic model_dump . **#92** RUFF lane
(+194 fixes; DELETED
273 never-compiled scripts/generate_khelacharya_flyer.py ? line 39 was a literal AI-
placeholder sentence) . **#93** MYPY
274 lane (7-module quarantine w/ ratchet rule) . **#94** SQLite conn-leak fix (340?4
ResourceWarnings) . **#95** coverage
275 floor 70 (CI-measured 72%) . **#96** config unknown-key warning . **#97** WSL
tilde/abspath . (**#98** Node-24
276 actions bump = owner-run chip session). Suite **373 green**; CI = 4 lanes
(smoke/lint/mypy/full+cov-floor).
277 - NEW session gotchas: `Select-Object -Last N` on a background pipe loses EVERYTHING if
the process is killed ? pipe
278 long background runs through `Out-File` instead; **sibling sessions flip branches in
THIS shared checkout** (the
279 node24 chip did, mid-session) ? re-verify `git branch --show-current` before every
commit; idle/suspend KILLS
280 background runs silently (smoke run #3 died mid-Gemini, no notification, stale report
on disk).
281 - ?? NEXT: unchanged owner-gated tracks (UI-alpha PR-2 upload/download per
`docs/UI_ALPHA_EXECUTION_PLAN.md`; corpus
282 growth; ship-gate target). Deferred-with-intent from the sweep: ruff I/UP rule groups;
yolo_hybrid P3/P4 fold-in;
283 pyproject packaging (AI-8 ? owner deselected); extra='forbid' flip (awaits gemini_*
alias retirement).
284
285 **Checkpoint:** 2026-06-11 (CI Node-24 action bump ? **#98 MERGED**, master `7717e15`).
GitHub forces actions onto
286 Node 24 from 2026-06-16 (Node 20 removed from runners 2026-09-16); ci.yml pinned
checkout@v4 + setup-python@v5 (Node 20).
287 Verified latest majors via `gh api releases/latest`: **checkout v6** (v6.0.3 ? NOT v5 as
the run-log warning suggested)
288 + **setup-python v6** (v6.2.0). Bumped both in all 4 lanes; all green on the PR (lint
11s / mypy 29s / smoke 15s /
289 full+cov 1m27s), squash-merged on verified rollup. PR #97 (fix/wsl-tilde-quoting) was
open + untouched throughout;
290 working tree restored to it (PR-2 doc edits intact).
291
292 **Checkpoint:** 2026-06-11 (8?9 QUALITY SWEEP IN PROGRESS ? owner approved 8 AIs via
question UI; AI-8 packaging NOT
293 selected). **MERGED on green CI: #90** (hygiene: .coveragerc repointed post-#40 +
TEST_RUN_SUMMARY ? docs/archives),
294 **#91** (pydantic .dict()?model_dump, 7 sites, warnings 15?0), **#92** (? RUFF LANE:
ruff.toml E/F/W/B py38 +
295 fix 194 violations + `lint` CI job; **DELETED scripts/generate_khelacharya_flyer.py ?
never compiled, line 39 was a
296 literal AI-placeholder sentence**, invisible because nothing swept scripts/), **#93**
(MYPY LANE: mypy.ini lenient
297 green baseline, explicit 7-module quarantine w/ ratchet rule ? main.py+5
segmenters+registry; fixed 18 annotation-level
298 errors), **#94** (? DB FIX surfaced by coverage: `with sqlite3.connect()` = transaction
scope NOT close ? all 24
299 Database call sites leaked connections; new `_connect()` CM; ResourceWarnings 340?4),
**#95** (COVERAGE FLOOR:

```

```

300    full-suite lane runs --cov, **--cov-fail-under=70** vs CI-measured 72% ? ratchet-only),
**#96** (config unknown-key
301    warning: dotted-path [CONFIG WARNING] for typo'd/dropped keys at all nesting levels;
repo config.json verified clean;
302    extra='forbid' still deferred). Suite 349? **351 green**. CI = 4 lanes
(smoke/lint/mypy/full+cov). ?? NEXT: **AI-7
303    hybrid-twin collapse** (shared trajectory-hybrid base, fix shot_count drift, dedupe
_wsl_bash, equivalence tests) +
304    **post-AI-7 quality-eval checkpoint** (owner: verify ~0.66-class evals stick; GPU free,
Gemini key ROTATED, live calls
305    OK). Deferred-with-intent this sweep: ruff I/UP rule groups; yolo_hybrid P3/P4 fold-in
(not as follow-up).
306
307    **Checkpoint:** 2026-06-11 (PR #88 ? **#89 MERGED** ? KhelSutra UI-alpha docs bundle on
master `34cf26c`). #88
308    conflicted (branch cut from `feat/async-job-model` pre-squash of #87) ? per
`scratch/PR88_FIX.md`, recreated as
309    `docs/khelsutra-plan-2` off post-#87 master keeping 6 paths (UI_PROPOSAL, HOSTING_PLAN,
UI_ALPHA_EXECUTION_PLAN,
310    mockups/, DECISIONS **ADR-009**, NEXT_STEPS sync); #88 closed + branch deleted. CI
green-gated (statusCheckRollup),
311    squash-merged.
312    - **Dropped from #89, preserved for PR-2 as UNCOMMITTED working-tree edits on master:**
(1)
313    `docs/DEFERRED_HARDENING_PLAN.md` "#16 DONE ? MERGED (PR #87)" status hunk (also saved
as
314    `scratch/PR2_deferred_hardenening_status.patch`); (2) `docs/UI_ALPHA_EXECUTION_PLAN.md`
working-agreement rewrite
315    (Cowork designs + leaves `scratch/PRn_HANDOFF.md`; **Claude Code executes all git/PR
mechanics** ? ratified
316    2026-06-11). Don't `checkout --` these away; they commit with PR-2.
317    - ?? Next UI-alpha slice = **PR-2 (upload/download)** per
`docs/UI_ALPHA_EXECUTION_PLAN.md`.
318    - ? **SLATE CLEANED (2026-06-11):** indexer = NO open PRs, master==origin (`34cf26c`),
**all 23 stale local branches
319    pruned** (every one mapped to a MERGED PR #62?#87 before `-D`); only `master` remains.
Collector = NO open PRs, on
320    master. Intentional carryovers (NOT mess): the 2 uncommitted PR-2 doc edits (above),
untracked `scratch/` handoffs,
321    untracked `artifacts/rally_tal_gen0/0.1.0/` (Gen-0 bundle, ADR-008). ?? Collector has
a stray uncommitted
322    `.gitignore` edit (adds `artifacts/` ignore w/ ADR-008 comment) ? mega-session
leftover, owner to commit (micro-PR)
323    or fold into the next collector PR.
324
325    **Checkpoint:** 2026-06-10 (GEMINI BATTERY COMPLETE ? full coverage; owner's billing was
fine, postpay active; ALL
326    failures were burst/capacity artifacts). Merged this stretch: **#80** (chunker re-
encode/downscale, from the chip
327    session ? reviewed+merged), **#81** (rescued the chip branch's orphaned run_signals test
via cherry-pick), **#83**
328    (provider: 6 generate attempts exponential backoff+jitter ? saved chunks live), **#84**
(tracked
329    `eval_baselines/gemini_wrapper_2026-06-10.json`). Suite 336 green.
330    - **FINAL QUALITY TABLE (IoU .5 vs human golden labels):**
331    - **Mahadevpura** (clean-ish, 31 rallies): heuristic 0.657 · Gen-0 0.744 · **two-stage
WASB+Gemini refine 0.83**
332    (P .81 R .84, 32 preds) · **full-video wrapper 0.93** (P 1.0 R .87).
333    - **Kushagra** (multi-court, 32 rallies, FULL 7/7 coverage + audio): heuristic 0.157 ·
Gen-0 0.561 ·
334    **wrapper 0.66** (P .69 R .62, ±1.2s; FPs = background-game pickups + one 53s dead-
time merge ? exactly the
335    contrast-metric confound).
336    - **STRATEGY (refined):** clean footage commoditized (0.93 cheap); **multi-court drops
the wrapper to 0.66 = the

```

```

337     owned model's ship target** (Gen-0 is 0.10 behind at n=6 ? within corpus-growth
reach). Two-stage refine quality
338     (0.83) < full wrapper on short videos; its value = COST on long academy tapes (only
candidate clips uploaded).
339     Wrapper reliability is genuinely flaky (503 waves all session) ? provider-fallback
architecture is load-bearing.
340     - Artifacts: `output/mahadevpura_gemini_refined.csv`,
`output/MahadevpuraSingles_highlights.mp4` (demo reel),
341     Kushagra preds under the ORIGINAL corpus MD5 now. Total Gemini spend today: ~$1-2.
342     - PENDING/owner: rotate the chat-pasted key; ship-gate F1@tIoU decision (multi-court
target now has a number: 0.66);
343     C13 interviews + camera pilot; repo rename.
344
345     **Checkpoint:** 2026-06-10 (GEMINI BASELINE LANDED ? owner's "gradual ramp-up" theory
CORRECT). New key in gitignored
346     `.env` (rotate later ? pasted in chat). The "credits depleted" 429 was a BURST-THROTTLE
artifact: text ping + serial
347     upload worked ? **PR #79 merged**: gemini segmenter now honors `indexing.ai_max_workers`
(was hardcoded 5; config ships
348     1=serial ? at 5 workers the same video 429'd, at 1 it worked).
349     - ? **WRAPPER BASELINE (gemini-flash-latest, ~$0.2-0.5 total): Mahadevpura F1 0.93** (P
1.00 R 0.87, mIoU 0.80, ±0.7s;
350     3 of 4 misses = short rallies Gemini FOUND but our min_segment_duration=3.0 rejected)
vs Gen-0 learned 0.744 vs
351     heuristic 0.657. **Clean footage = commoditized.**
352     - ?? **MOAT TEST (Kushagra multi-court, windowed n=12 GT): wrapper F1 0.56** (P .54 R
.58 mIoU .69) ? Gen-0 learned
353     0.561, both >> heuristic 0.157. **The multi-court confound hurts Gemini too ? the
wedge survives on the target
354     footage class.** Caveats: only chunks 2+4 answered (5/7 lost to GENUINE gemini-flash-
latest capacity 503s ? retry
355     off-peak for full coverage); visual-only (audio stripped in downscale); 1280p re-
encode (4K chunks were 1.3GB >
356     500MB upload cap ? the audit's silent-oversize-skip fired live; fix = chip
task_44e72c41 ? see PR #80 checkpoint
357     below). Kushagra preds keyed under downscaled-file MD5 (`output/kushagra_1280.mp4`).
358     - ? **DEMO ARTIFACT: `output/MahadevpuraSingles_highlights.mp4`** ? real 27-rally reel
compiled by VideoCompiler
359     (the once-broken module) from Gemini segments. The academy-pitch object exists.
360     - **STRATEGY READ:** product path = Gemini for clean footage TODAY (cheap, F1 0.93) +
owned model's job = multi-court
361     (0.56 is the bar to beat there, not 0.93); 503 flakiness = real reliability cost of
the wrapper (provider-fallback
362     story strengthens). PENDING: gemini_refine two-stage (503-gated, run off-peak), full
Kushagra re-run, PR #80 merge.
363
364     **Checkpoint:** 2026-06-10 (4K gemini oversize fix ? **MERGED: #80 + #81 + #82; master
`f01387e`, CI rollup SUCCESS**).
365     Root cause of the live "0 segments / success" on KushagraYashAviFirstVideo.MP4 (4K):
gemini segmenter stream-copies 90s
366     chunks ? all 7 at 832?1299MB > provider `MAX_UPLOAD_MB=500` ? every chunk silently
skipped. Fix (3 layers): (1) chunks
367     RE-ENCODED + downscaled to new declared knob `indexing.ai_chunk_scale_width` (default
1280 = gemini_refine's
368     `--clip-scale-width`; 0 = legacy stream-copy; also fixes keyframe drift); (2) skip is
LOUD ? provider
369     `metrics["oversize_skipped"]` + segmenter aggregation (chunked + single-file) ?
`logger.error` + NEW thread-safe
370     `backend/reporting/run_signals` channel (incr/pop by video_id; popped unconditionally in
`emit_indexer_report`);
371     (3) report flags it ? `ai_handoff.details.oversize_clips_skipped` + spec.yaml ERROR rule
`oversize_clip_skipped`.
372     Suite 325?336 green** locally (CI runs fewer: reporting tests importorskip obsreport,
absent there BY DESIGN ?
373     `tests/test_run_signals.py` deliberately has NO importorskip = the CI-visible half of

```

```

the channel). **Merge dance
374 (process note):** owner merged #80 mid-session (auto-deleting the remote branch) WHILE
the follow-up test commit was
375 being pushed ? push re-created the branch orphaned; owner rescued it as #81, 39s before
my cherry-picked #82 ?
376 **#82's squash `f01387e` = EMPTY commit** (harmless dup in history). Lesson: re-check
`gh pr view` state before pushing
377 more commits to a PR branch ? owner + sibling sessions work concurrently (this
checkpoint file too: another session
378 prepended the GEMINI-BASELINE checkpoint mid-edit). ? **OWNER/SIBLING WIP IN TREE
(uncommitted, do NOT clobber):**
379 `backend/providers/gemini.py` ? generate retries 3?6 w/ exponential backoff + jitter
(~3?96s) for the 503 capacity
380 waves. ?? commit/PR the 503-backoff WIP . machine-side 4K re-run (chunks should be ~tens
of MB) . task-17 battery
381 still pending Gemini billing top-up.
382
383 **Checkpoint:** 2026-06-10 (M0 GREENLIT SEQUENCE EXECUTED ? session wrap). Owner
ratified ADR-008 ("split driven by
384 productionization, not isolation") + greenlit M0.4+M0.1. **All 5 steps shipped, every PR
merged on verified-green CI:**
385 - **#73 ADR-008** (DECISIONS.md): training in-repo behind an executable wall; artifact-
bundle boundary; ship-gate
386 product-side; pre-committed split trigger ? `sports-temporal-models` at Gen-2/second-
consumer.
387 - **#74 canonical GT loader** `backend/eval/gt_loader.py::load_gt_intervals` ? closes
the two-loader fork (accepts BOTH
388 `start,end` and `start_time,end_time`); legacy loaders delegate; round-trip contract
test; LOVO 0.626 re-verified.
389 - **#75+#76 featurizer promotion + import wall**:`backend/features/rally_seq.py`
(FEATURE_SCHEMA v1, fps semantics
390 explicit) imported by train AND serve; `training/` skeleton (5 wall rules,
requirements-train.txt); smoke tests:
391 runtime no-training-in-backend.main, static sweep, featurizer purity (torch/cv2
blocked; importorskip numpy on the
392 minimal CI lane ? #76 hotfixed a red lane after #75 merged early; merges now gated on
verified SUCCESS rollout).
393 - **collector #19 ? n=6 OWNED CORPUS IN THE SoR**`all 6 golden matches via `import-
local --normalize` (262 rallies,
394 0 dropped, Proprietary, explicit fps); commercial export now 6 real matches + demo
entry.
395 - **#77 Gen-0 harness** `python -m training.gen0.harness --manifest ? --floor
eval_baselines/heuristic_lovo_n6.json
396 --version 0.1.0`: train?honest LOVO?ship-gate?artifact bundle (gitignored weights.json
+ manifest.json w/ schema pin,
397 calibration, lineage gt_hashes, attestations incl. C3, gate verdict). **First real
bundle built: LOVO 0.626, floor
398 PASSED (0.480), paired sign test 5/6 wins p=0.109 = honestly NOT significant,
ship=False** (blockers: target
399 undefined, C3, significance). Committed floor baseline
`eval_baselines/heuristic_lovo_n6.json`. Suite **325 green**.
400 - ? **Gemini baseline STILL BLOCKED after owner provided a NEW key (2026-06-10 later):**
key written to gitignored
401 `.env`, validated via `models.list` (55 models incl. `gemini-flash-latest`; **config
migrated off deprecating
402 gemini-2.5-flash ? `gemini-flash-latest` alias, PR #78 merged**`gemini_refine already
defaulted to it). But inference
403 = 429 RESOURCE_EXHAUSTED "prepayment credits depleted" ? **the AI-Studio PROJECT
behind the key has no credits** (key
404 auth ? billing). Owner: add credits at https://ai.studio/projects (+ consider rotating
the chat-pasted key). Planned
405 battery on unblock (in task 17): wrapper Mahadevpura ? score; wrapper Kushagra (multi-
court MOAT test); two-stage
406 gemini_refine; compile a real highlight reel (academy-demo artifact). Pipeline handled
both failed runs gracefully

```


407 (validators pass, reports emitted, artifacts cleaned); shared-genai-client thread-safety finding observed live.

408 - **C13 field plan (owner executing):** recruit interviewer; camera idea RATIFIED w/ 3 constraints (consent-at-capture/

409 adults-only for training; demo via Gemini-or-human-polish, NOT gated on owned model; interviews?cameras sequencing).

410 - ?? NEXT: owner tops up Gemini billing (wrapper baseline) · owner sets the ship-gate F1@tIoU target · grow corpus

411 (camera pilot videos ? import-local ? re-run harness) · repo rename · M0.1 event-index contract in collector.

412

413 **Checkpoint:** 2026-06-10 (Roora PB+TT pilot review ? REPORT DELIVERED to owner; owner emails Roora himself) ·

414 ? Owner cross-checked the rally tables in VLC ? **CONFIRMED: every boundary off ?0.5s (up to ~2s)** vs the ±0.3s bar

415 (Brief §4). **Root cause (my forensics): all 24 boundaries sit EXACTLY on the 30-frame sampling grid** (~0.5s at

416 59.94fps) ? first/last box never on the true contact/dead frame ? structurally cannot meet ±0.3s. Also decoded their

417 bogus `start/end time` attrs = **frame÷900 (decimal minutes @ assumed 15fps, ~4× off wall-clock)**.

418 Report findings: boundary accuracy (CRITICAL), skip-rallies-in-progress-at-file-start/end convention (PB rally#1 starts at frame 0 ?

419 shouldn't have been labeled), dead-time+warmup reminders for Phase-B, schema cut to exactly

420 `rally\_number,shot\_count,ending\_reason,last\_shot\_outcome,note` (drop their time attrs + scores), field quality

421 (PB outcome='0' off-vocab; TT 5/6 'other' + missing required notes; TT#5 intra-rally other|winner). Per owner:

422 NO redelivery demand in the report (he words the email himself; thread = Jamshad@roora on avidullu@live.com Outlook).

423 **Artifacts:** PDF `Annotation Setup\6 June 2026\Foodball\Roora Pilot Review - Pickleball and Table Tennis

424 (2026-06-10).pdf` + md record at `Annotation Setup\Roora Pilot Review ? Pickleball & Table Tennis (2026-06-10).md`

425 (generator: `%TEMP%\roora\make\_review\_pdf.py`). ? Local guideline docs are **DRAFT v2** ? owner mentioned "v7";

426 no v7 found on disk (possibly in the email thread; report cites v2). Spawned chip `task\_c55f0be6`: collector

427 `rally\_cvat.py` ATTR_MAP drops `shot\_count` (present in Roora XML, empty in our CSVs) ? OUR bug, kept out of the

428 vendor report. ?? owner sends email; Roora's badminton long-video + redelivery decisions pending owner.

429

430 **Checkpoint:** 2026-06-10 (REPO-TOPOLGY decision panel + C13 plan ? AWAITING OWNER RATIFICATION). Owner leaning

431 M0.4+M0.1 greenlight; asked: model in-repo vs separate repo consumed as a segmenter (owner prefers isolation). Ran a

432 4-agent judge panel (`wf\_471a9cdd-d1d`; full output `?\tasks\wnksv545g.output`).

433 **Verdict (my rec): "artifact boundary NOW, repo split at Gen-2" ? M0.4 lands as a top-level `training/` tree in the indexer w/ executable isolation

434 (import-smoke asserts backend.main loads no training module; own requirements-train.txt + 3rd CI lane), featurizer

435 promoted to `backend/features/rally\_seq.py` (single source, train==serve), artifact bundle = gitignored weights

436 (private Releases) + committed manifest (feature-schema ver, normalization, calibration, corpus ref+gt_hash,

437 consent/license attestation incl. C3), ship-gate stays product-side (regression.py extension); **pre-committed split

438 trigger** to a named repo (`sports-temporal-models`) when Gen-2/cloud-training or a 2nd artifact consumer arrives.

439 Decisive evidence: Gen-0's input = the indexer's OWN featurization (mid-pipeline seam ? WASB/Gemini video-boundary);

440 indexer not pip-installable ? split today = copy-paste = the observed drift class;

```

obsreport "precedent" never actually
441     exercised (pin is a commented line, repo unpublished, CI skips). **FIRST contract to
freeze: golden-label/corpus
442     contract** ? indexer carries TWO incompatible GT loaders (annotations.py `start,end` vs
calibrate_wasb `start_time,
443     end_time` ? fed to 7 modules) vs collector export `start,end`; consolidate + round-trip
test BEFORE any training epoch.
444     Panel extras: n=6 corpus's durable home = loose `Annotation Setup\` folder (neither
repo!) ? M0.1 first deliverable =
445     hash-addressed in collector; training fixtures synthetic-only until consent ledger
exists; run the Gemini-wrapper
446     baseline before/with M0.4. **C13 field plan:** owner recruiting an interviewer; my rec =
bake in the camera idea
447     (paid recordings = consented owned multi-court corpus + demo), demo served via Gemini
path or human-polish (owned model
448     honest 0.626 ? >0.8 yet; the comparison doubles as the wrapper baseline). ?? awaiting
owner ratification of topology +
449     M0.4/M0.1 scope; then: ADR + GT-loader consolidation + featurizer promotion + harness
scaffold.
450
451     **Checkpoint:** 2026-06-10 (CORPUS RECLASS EXECUTED + NEXT_STEPS refreshed ? session
wrap). Owner approved + I ran the
452     corpus correction: `reclassify-licenses --apply` (23 ShuttleSet broadcast rows ?
annotation_license=MIT, video=Unknown,
453     is_commercial=False; **zero annotations deleted**, 1,917 rallies intact) + `curate
export` (commercial manifest now
454     owned-data-only: 1 video/2 segments; stale ShuttleSet CSVs removed from the gitignored
eval_export; eval-only re-export
455     available via `--include-staging --include-noncommercial`). **Merged: collector #17**
(DB + .gitignore backup rule);
456     pre-change backup at `data/sports_metadata.backup-2026-06-10.db` (gitignored). Quality
signals UNAFFECTED (n=6 corpus
457     never flows through the collector DB; learned 0.626 re-verified earlier). **Indexer #72
merged:** NEXT_STEPS.md
458     refreshed (sweep banner, 0.626 vs 0.480, M0.3 partially-done with corrected
distill_local pointer, NEW prioritized PMF-
459     validation item, durable artifact paths). **Owner is now reading NEXT_STEPS to chalk out
next steps** ? open owner
460     decisions stand: greenlight M0.x (rec M0.4+M0.1), confirm collector=SoR, ship-gate
F1@tIoU, repo rename, C13 interviews
461     / Gemini-wrapper baseline. Both repos clean on master, 0 open PRs.
462
463     **Checkpoint:** 2026-06-10 (HIGH-RISK REMEDIATION COMPLETE ? owner approved "everything
A?K, open+merge PRs after CI

```

76. assistant (2026-06-13T19:35:11.356Z)

This is the append-only live checkpoint log. I'll prepend a concise new checkpoint for PR #116 (a meaningful checkpoint per the working agreement).

77. user (2026-06-13T19:35:26.920Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

78. user (2026-06-13T19:35:33.959Z)

```

1     ---
2     name: current-state
3     description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4     metadata:
5         node_type: memory
6         type: project
7         originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301

```

```

8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-13 (RALLY-QUALITY TRACK: HARNESS + PERSON-CUE EXTRACTION SHIPPED
? **PR #114 + #115 MERGED**,
13     master `df7d5df`; suite 421?441 green; 0 open PRs). **PR #115 (fusion P1):** extracted
the torchvision person detector
14     + ByteTrack + velocity windowing out of `yolo_hybrid` into
`backend/pipeline/detectors/person_detector.py`

```

79. user (2026-06-13T19:35:51.407Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

80. assistant (2026-06-13T19:35:36.698Z)

Now retrying the checkpoint prepend.

81. assistant (2026-06-13T19:36:13.965Z)

Done. Here's the summary.

What I delivered

PR [#116](https://github.com/avidullu/badminton-highlight-indexer/pull/116) ? branch `docs/analyzers-and-reconcilers` off `master`, **docs-only, owner-gated, not merged**. Contains exactly two files:

- **docs/ANALYZERS_AND_RECONCILERS.md** (new, 367 lines ? peer-sized to the fusion plan)
- A one-line cross-link in **docs/README.md**

The doc

The **spine** is the owner directive, threaded through every section: **analyzers/reconcilers emit signals* (value + confidence); the **scoring model** learns to weight them* ? no *`if/else`* in the decision path.* I anchored it to a pattern that already exists: *`experiment.predict()`/_gate_mask`* already treat *`net_crossings`/\_players\_in\_court`* as feature columns.

- **Seam 1** ? injectable *`WindowAnalyzer`s (forward):* a registry mirroring *`segmenter\_registry`*, running pure per-window cues over the already-produced per-frame-aligned signals (WASB trajectory + *`PersonDetector.detections\_per\_frame`* + *`rally\_gate`* net). *`AnalyzerResult`* carries all three roles ? confidence-scored features ? *`candidate\_segments.stats`*, events ? the *`events`* table, optional boundary hints. Worked example: the **service-fault** analyzer (shuttle halts in the net region ? *`service\_fault`* signal + *`net\_hit`* event + END hint). Injects at the FusionSegmenter *`\_enrich\_candidate\_stats`* hook; flag-gated, default OFF.
- **Seam 2** ? reconciliation pass (backward): a report-first "linter for rally decisions" with four worked *`ConsistencyRule`s* (trailing-dead-time ? truncate *(the owner's misfire fix)*, no-net-crossing ? drop, over-merged ? split, boundary-slack ? tighten). **Report-first**, never silent; expressible as a harness *`Variant`* transform so its lift is LOVO-measurable on the 6 golden videos **before** auto-apply. Idempotent, fixed precedence.

Plus: the unifying pipeline diagram, JSON-serializable contracts, persistence (stats/events/corrections via the *`user\_version`* migration ladder), the no-regression discipline, composition with fusion P2, trade-offs/risks, a phased A1?A3 / R1?R3 PR plan, and ADR-007 references.

One thing worth flagging

The task warned the fusion-P2 work might not be on master ? it isn't. **The entire fusion-P2 PR is sitting uncommitted in the shared working tree** (``fusion_hybrid.py``, ``fusion_features.py``, ``FusionConfig``, ``person_detector.detections_per_frame``, ``QUALITY_ITERATIONS.md``, and pre-existing README edits). I:

- designed against ``MULTI_SIGNAL_FUSION_PLAN.md`` and the real enrichment-hook seam, and
- **kept that WIP byte-intact** ? I isolated my single README line with git plumbing (the tree is CRLF / ``autocrlf=true``) and restored the owner's README to its exact session-start bytes, so my PR carries only my line, not the in-flight changes.

Because the doc cross-links ``QUALITY_ITERATIONS.md`` and the fusion hook (neither on master yet), the PR body notes: **merge #116 after, or together with, fusion-P2.**

I also refreshed the live ``current-state.md`` memory checkpoint. Implementation stays owner-gated ? nothing started.

82. user (2026-06-13T19:36:16.350Z)

can you pause and not merge the PR

83. assistant (2026-06-13T19:36:22.762Z)

I haven't merged it ? and I won't. PR [#116](https://github.com/avidullu/badminton-highlight-indexer/pull/116) is **open and untouched**, exactly as the task required ("do NOT merge ? owner reviews").

Current state:

- Branch ``docs/analyzers-and-reconcilers`` pushed, PR open, **no merge performed**.
- Nothing further is running ? I'm paused.
- The in-flight fusion-P2 working-tree WIP remains byte-intact.

It's all yours to review whenever you're ready. I won't take any further action on it unless you ask.

84. user (2026-06-13T19:37:09.103Z)

One more idea to add or push back is that we can add a per frame "shuttle in air or shuttle hitting racket or shuttle inactive! So that way if we can add a new vision detector which tells the state of shuttle we could add that more signal but in this case we need a layer of frame processing like wasb

85. assistant (2026-06-13T19:38:36.533Z)

Good idea ? and you've actually put your finger on the exact thing that makes it **different** from the two seams in the PR. Let me place it correctly and push back on one part.

It's a detector, not an analyzer or reconciler

The shuttle-state classifier (``in_air`` / ``racket_contact`` / ``inactive``) **processes raw frames ? a new per-frame signal**, exactly like WASB. That puts it in the **detection layer**, alongside WASB / PersonDetector / the audio hit-detector ? i.e. it's a new **RallyCue` producer** in ``MULTI_SIGNAL_FUSION_PLAN.md`` §2, **not** a ``WindowAnalyzer``.

The distinction is the whole spine of the quality track (EXPERIMENT_HARNESS §2): **expensive DETECTION** (GPU/WSL, frame-processing, risky) vs **cheap DECISION** (pure functions over already-produced signals). My analyzers/reconcilers are the cheap-decision layer ? they **consume** signals; they don't produce new ones. Your idea lives one layer up, in the expensive layer. They compose, but the cost/risk/labels story is completely different, which is why it belongs in the fusion doc, not the analyzers doc.

How it composes once it exists (this is the nice part):

- An **analyzer** turns the persisted per-frame state into window features ? ``racket_contact_count`` (? shot cadence), ``inactive_ratio`` (? dead-time) ? which flow into the scorer as columns. No special-casing.
- The **reconcilers** get **stronger evidence**: the ``trailing_dead_time`` and ``over_merged`` rules

currently infer inactivity from a velocity threshold; a learned `inactive` state is a more direct, occlusion-robust evidence source for "truncate here / split here."

The push back: build the derived version first

Here's the honest problem with jumping straight to a new vision model:

1. **No training labels.** Golden labels are rally `[start,end]` times ? we have **zero** per-frame shuttle-state labels. So you'd either run a new (expensive, minors-gated) per-frame annotation task, or weak-label from trajectory+audio. But if you can weak-label state from existing signals, **you can just compute the feature from those signals directly** ? no model needed.
2. **Most of the state is already derivable, for free:**
 - `in\_air` vs `inactive` ? WASB velocity threshold + visibility (already computed).
 - `racket\_contact` ? **direction reversals** of the trajectory (already noted as `direction\_reversals` in fusion §4 ? "each reversal ? a shot") **and** the **audio thwack** ? which is literally "shuttle hitting racket" and is **already scaffolded** (`backend/pipeline/audio/hit_detector.py`, ADR-007's #1 cue).
3. **A second WSL frame pass ? doubles GPU cost** (separate process from WASB, no shared backbone).

So a **`shuttle\_state` analyzer** over signals we already have (velocity + visibility + direction-reversals + audio) is a free v0 ? persist it, A/B it on the 6 golden videos, and see if it lifts. **The dedicated per-frame vision model only earns its keep where the derived version hits a ceiling:** WASB trajectory gaps ? occlusion, motion blur, fast smashes ? where there's no `(x,y)` to threshold but an appearance model could still call `racket\_contact`. That's its unique, real value, and it's the right trigger to build it (gated on labels existing, same blocker as feature-level fusion).

My recommendation

Stage it: (a) **`shuttle\_state` analyzer** from existing signals now ? measure; (b) promote a learned **per-frame detector** only if (a) is blocked by trajectory gaps and the labels exist. That keeps it on the repo's "persist once, re-score, promote on lift, don't build capability before measuring" rails.

Want me to fold this in? I'd add a short subsection to **`MULTI\_SIGNAL\_FUSION\_PLAN.md`** (the shuttle-state **producer** ? derived-cue-first, vision-model-when-trajectory-gaps-bite, label blocker) and a one-paragraph **"consuming a richer detector"** note in the analyzers doc showing how it slots into both seams. That's a new commit on the open PR branch ? **I won't merge**, and I'll keep the in-flight fusion-P2 WIP byte-intact as before. Or I can leave the docs as-is and just keep this in the thread. Your call.

86. user (2026-06-13T19:38:38.847Z)

one more is the "remove practice and warmup" processor, all its job is to identity start and end of warmup so that the scorer may consume it

87. assistant (2026-06-13T19:39:59.792Z)

Right, and this one has a structural wrinkle worth making explicit: a warmup detector is a **different scope** from either seam ? and it's the single best illustration of why the "emit a signal, don't branch" spine matters.

It's a session-scope producer, not per-window or per-segment

- A **`WindowAnalyzer`** sees **one window**.
- A **`ConsistencyRule`** sees **one predicted segment** vs its evidence.
- Warmup is a **span** over the whole timeline ? it covers many windows, and you can only find it by looking at the **sequence** of activity, not any single window. So it's a new sibling tier: a **`SessionAnalyzer`** (runs once per video over all the persisted window/frame signals) that

emits a `warmup` span `[start, end]` + a per-window membership feature.

That's a small, clean contract extension: `analyze(all\_windows, signals) -> { spans, per\_window\_features }`. And note it needs **no new frame detector** ? unlike the shuttle-state idea, it's pure decision-layer math over signals we already persist, so it's fully offline/re-scorable.

You framed the consumption exactly right ? and that's the whole point

```
> "so that the scorer may consume it"
```

Yes. It emits `in\_warmup` (overlap fraction, 0?1) as a **feature column**; the scorer **learns** to down-weight windows inside the span. **Warmup** is the poster child for no-special-casing: the naive version is `if t < warmup\_end: drop` ? a hand-coded branch hanging off a brittle, fuzzy boundary. If warmup-end is mis-estimated by 20 seconds, a hard cut **deletes** a real opening rally. As a weighted feature, a wrong boundary just slightly mis-weights a couple of windows the other cues can rescue. Same reason the reconcilers are report-first.

It's measurable today ? no new labels needed

The golden labels already treat warm-up knock-ups / pre-serve as **dead time = negatives** (Roora starts a rally at serve-contact; fusion \$1 lists "warm-up knock-ups" as a named precision target). So you don't need an explicit "this span is warmup" label ? an `in\_warmup` feature is directly A/B-able on the 6 golden videos: does it lift precision? That's the promotion gate.

The honest push back

1. **Warmup** ? a long rally, structurally. Without a scoreboard or serve-OCR, the discriminators are weak individually ? continuous activity, low inter-window dead-time, no serve-reset cadence, position-at-start. This is **exactly** the temporal confound the repo already documented as the hard case: QUALITY_ITERATIONS' standing lesson says the multi-court dead-time confusion is "rally-vs-non-rally activity in the same place and frequency band... only the learned model's temporal features crack it." Warmup-vs-play is the same shape ? which is itself the argument for **feature-the-model-weights**, not heuristic span-cutting.
2. **"Practice"** isn't only at the start. Between-games knock-ups, a coach feeding drills mid-session ? so don't bake in a hard position prior beyond a v0. Start with "leading contiguous span" as the cheap v0, generalize later.
3. **Nice synergy** with your other two ideas: the cleanest warmup discriminator is **cadence** ? warmup is continuous hits **without** the point-ending pause rhythm. That's precisely what the **shuttle-state** classifier (`racket\_contact` cadence) and the scaffolded **audio thwack** detector produce. So warmup detection is a natural **consumer** of those signals, not a separate model.

Recommendation

Model it as a **SessionAnalyzer** emitting a `warmup`/`practice` span + an `in\_warmup` per-window feature; default OFF; promote on measured golden-precision lift. v0 = leading-span heuristic over existing window timings (free); v1 = learned, fed by shuttle-state/audio cadence once those exist. Reconcilers can **also** use the span as evidence to down-weight a whole cluster, but the primary path is the feature ? scorer.

I now have **two new ideas** to fold in (shuttle-state producer + this session-scope warmup analyzer). Want me to make **one commit** to the open PR that:

- adds the `SessionAnalyzer` tier + the warmup worked-example to `ANALYZERS\_AND\_RECONCILERS.md`, and
- adds the shuttle-state **producer** subsection to `MULTI\_SIGNAL\_FUSION\_PLAN.md` with a consuming note in the analyzers doc?

Still **won't merge**, and I'll keep the in-flight fusion-P2 WIP byte-intact. Or hold both in-thread until you've collected all your ideas ? say the word.

88. user (2026-06-13T19:40:01.699Z)

please fold these as well and extend your doc

Also I hope you are treating this as a framework being built not specifically for this project but also as a research paper. Detecting rallies is a reasonably hard problem

89. assistant (2026-06-13T19:41:58.069Z)

Folding both in ? and you're right to push the altitude. I'll reframe the doc around a **general signal-fusion framework** (producers at multiple temporal scopes ? a scorer that arbitrates ? an auditable correction layer), with rally detection as the motivating hard problem, and add a positioning section for the research framing. I'll keep this self-contained in **my** doc (the shuttle-state **producer** formally belongs in the fusion plan, which the in-flight fusion-P2 PR owns ? I'll document it here and cross-reference rather than edit that file concurrently).

Let me rewrite the doc.

90. user (2026-06-13T19:41:59.229Z)

```
1      # Analyzers & reconcilers ? two extensibility seams for rally-detection quality
2
3      **Status:** **DESIGN ? owner-gated, NOT started.** A peer to
4      [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) and
5      [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md).
6      This doc specs two *seams* ? a forward one (injectable window analyzers) and a backward
7      one (a
8      reconciliation pass) ? that let us keep piling on clever rally signals **without the
9      pipeline becoming a
10     pile of special cases.** Docs-only; no implementation code. Each phase below is a
11     separate, owner-gated PR.
12
13     > **Companion docs:** the fusion seam these plug into is
14     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md)
15     > (`RallyCue` = swappable producer; the `_enrich_candidate_stats` hook; the
16     scale/translation-invariant
```

91. user (2026-06-13T19:44:40.329Z)

<tool_use_error>File has been modified since read, either by the user or by a linter. Read it again before attempting to write it.</tool_use_error>

92. assistant (2026-06-13T19:42:07.462Z)

Approved ? folding both ideas in and updating the PR (no merge). Writing the extended, framework-framed doc now.

93. assistant (2026-06-13T19:44:40.325Z)

A signal-fusion framework for rally detection ? analyzers, scopes & reconcilers

****Status:**** ****DESIGN ? owner-gated, NOT started.**** A peer to
EXPERIMENT_HARNESS.md and
MULTI_SIGNAL_FUSION_PLAN.md.
This doc specs a small **framework** for evolving rally-detection quality: producers that emit ****signals****
at three temporal ****scopes**** (frame / window / session), a ****scorer**** that **learns to weight**
them, and an
auditable ****reconciliation**** pass that lints the decision against the evidence. The goal is to
keep piling
on clever cues ****without the pipeline becoming a pile of special cases.**** Docs-only; no
implementation
code. Each phase below is a separate, owner-gated PR.

```

> **Companion docs:** the detection-layer producers (shuttle / person / audio / shuttle-state)
live in
> [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md) (`RallyCue` = swappable producer;
the
> `_enrich_candidate_stats` hook; scale/translation-invariant feature discipline, §3.2); the no-
regression
> A/B + LOVO measurement these ride is [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md)
(`Variant`, `compare`,
> `compare_unlabeled`); the scorer is the distilled logistic regression (**ADR-004**,
> `backend/eval/classifier.py`); the modular rally layer is **ADR-007**
> ([archives/decisions/DECISIONS.md](archives/decisions/DECISIONS.md)); the 2-layer mental model
is
> [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md); the live quality log is
> [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md).

> **Framing.** Rally detection is a genuinely hard temporal problem ? weak, multi-modal cues;
interval-level
> ground truth that is scarce and expensive; a confound (warm-up / knock-ups / multi-court /
held-shuttle)
> that looks identical to real play in any single frame or frequency band. This doc is written
as a
> **reusable framework**, not a badminton-specific patch: the contracts are sport-agnostic and
the design
> is meant to generalize to other weakly-supervised temporal-segmentation problems (§10). Where
that
> ambition touches research claims, §10 says plainly what is and isn't yet substantiated.

```

1. The spine ? NO special-casing in the decision path

This is the one principle everything else serves:

```

> **Every producer (analyzer, detector, reconciler) emits a SIGNAL carrying a value + a
confidence/presence.
> The SCORING MODEL consumes those signals and *learns to weight them*. We do NOT add hand-coded
`if/else`
> branches into the decision.**

```

****Producers propose; the scorer arbitrates; the experiment harness + golden set decide what gets promoted.****

That separation is what lets us add a net-hit service fault, a let-cord, a double-bounce, a shuttle-out

cue, a **warm-up span**, a **shuttle-state** ? each as *"register a producer + a feature column,"* never *"add*

a branch."**** A signal that doesn't earn its keep stays at weight ? 0; it costs nothing because it is an unused column, not a live `if`.

The pattern already exists in miniature: `experiment.predict()`

(`backend/eval/experiment.py:203`) treats

`net\_crossings` and `players\_in\_court` as ****feature columns**** ? `\_gate\_mask` is the trivial scorer

(threshold a column), `LogisticRegression` (ADR-004) the richer learned scorer (weight all columns), a

future scorer richer still. The decision path is **signals ? scorer**, with ****zero**** `if service\_fault`:

drop`. Analyzers feed it more columns; reconcilers run **outside** it as an audited post-pass.

...

DETECTION ? frame-scope producers (expensive, GPU/WSL, run ONCE)

WASB shuttle (x,y,v,vis) · person boxes+tracks · audio hits · shuttle-state(in_air/contact/inactive)*

? per-frame-aligned signals (+ net estimate, window bounds)


```

?
shuttle-seeded candidate windows (FusionSegmenter, P2)
?
ANALYZERS ? decision-layer producers (pure, offline, re-scorable) [forward seam]
?? window-scope (WindowAnalyzer): one candidate ? features + events + boundary hints e.g.
service-fault
?? session-scope (SessionAnalyzer): whole timeline ? labeled spans + per-window membership
e.g. warm-up
? every signal ? candidate_segments.stats as feature columns
?
SCORING / DECISION ? signals ? learned scorer (weights, NO branches)
_gate_mask (trivial) + LogisticRegression (ADR-004) + harness-pinned model
? predicted rally segments
?
RECONCILERS ? backward seam, "linter for rally decisions" (report-first, idempotent)
ConsistencyRule registry: decision-vs-evidence conflicts ? Issue + corrected set
?
stats · events · spans · corrections record · highlights EXPERIMENT HARNESS gates EVERY
promotion
(* shuttle-state = a richer frame detector; its PRODUCER lives in MULTI_SIGNAL_FUSION_PLAN,
consumed here ? §3.5)
...

---
```

2. The framework ? signals, scopes, and who arbitrates

The architecture has exactly three kinds of moving part. Capability is added by registering a producer; the decision logic never grows a branch.

****Producers emit signals at three temporal SCOPES.**** A signal at a coarser scope is computed by aggregating finer-scope signals; everything ultimately lands as a per-candidate feature column (or an event/span the scorer or a human consumes).

Scope	Producer	Sees	Emits	Cost	Examples
frame	detector (`RallyCue`)	raw frames / audio	per-frame signal	expensive	(GPU/WSL) WASB shuttle, person boxes, audio hits, shuttle-state (§3.5)
window	WindowAnalyzer` (§3.1)	one candidate window's aligned frame-signals	features (+conf), events, boundary hints	cheap, pure	service-fault, double-bounce, let-cord, shuttle-out
session	SessionAnalyzer` (§3.4)	the whole timeline of windows	labeled spans + per-window membership features	cheap, pure	warm-up / practice , game/break phases

****The scorer arbitrates.**** It is the only place a decision is made, and it makes it by **weighting columns**, never branching. Today: `_gate_mask` (threshold) + `LogisticRegression` (learned weights, ADR-004)`. The confidence column on every signal lets the scorer discount an uncertain producer instead of trusting a binary verdict ? the small-N-safe way to fold in a deterministic-ish cue.

****Reconcilers audit.**** After the scorer decides, the backward seam (§5) lints the decision against the persisted evidence and proposes corrections ? report-first, never silent.

****Four invariants make this safe to grow**** (the load-bearing design claims):

- Signal, not verdict** ? every producer's output is a weighted feature; the decision path has no producer-specific branch. ? unbounded cues without combinatorial special-casing.
- Expensive detection ? cheap decision** ? frame-scope detection runs once and is persisted;

```

window/
    session signals and the scorer are pure re-scoring (EXPERIMENT_HARNESS $2). ? experiments are
offline,
    deterministic, zero production risk.
3. **Report-first reconciliation** ? corrections are recorded and reviewable, never applied
silently. ?
    human-in-the-loop, auditable.
4. **Promote only on measured held-out lift** ? LOVO on the golden set + the `run_regression.py`
gate;
    default OFF; rollback = config flip. ? no-regression evolution.

---

```

3. The forward seam ? injectable analyzers

Pluggable producers ? registered like `segmenter\_registry`
(`backend/pipeline/segmenters/base.py:35`) ?
that run over the **per-frame-aligned signals already produced**: the WASB shuttle trajectory
(`.frame/.visible/.x/.y` + velocity), the person boxes per frame
(`PersonDetector.detections\_per\_frame`,
`person\_detector.py:256`), the net line (`rally\_gate.estimate\_play\_axis`, camera-agnostic), and
the window
bounds. Analyzers are **pure** ? no I/O, no torch/cv2 ? mirroring `fusion\_features.py`'s
discipline, so
their math is unit-testable in the GPU-free dev container.

3.1 Contracts (JSON-serializable; proposed `backend/pipeline/analyzers/`)

```

```python
@dataclass(frozen=True)
class Signal:
 # one scalar cue + how sure the producer is ? never a
 verdict
 value: float; confidence: float = 1.0 # persisted as `<producer>__<key>` (+
 `_confidence`), [0,1]

@dataclass(frozen=True)
class AnalyzerEvent:
 # ? the `events` table on confirmation ($4)
 type: str; timestamp: float; confidence: float = 1.0 # "net_hit"/"double_bounce"/? ; src-
 absolute s
 player: Optional[str] = None; details: Dict[str, Any] = field(default_factory=dict)

@dataclass(frozen=True)
class BoundaryHint:
 # OPTIONAL deterministic proposal ? never auto-applied
 kind: Literal["start", "end"]; timestamp: float; confidence: float; reason: str

@dataclass(frozen=True)
class AnalyzerResult:
 signals: Dict[str, Signal] = field(default_factory=dict)
 events: List[AnalyzerEvent] = field(default_factory=list)
 boundary_hints: List[BoundaryHint] = field(default_factory=list)

@dataclass
class WindowSignals:
 # read-only pre-produced bundle (no I/O inside the
 analyzer)
 start: float; end: float; fps: float; frame_width: float; frame_height: float
 shuttle: List["TrajectoryPoint"] # in-window points (.frame/.visible/.x/.y)
 persons: Dict[int, List[Tuple[int, BBox]]] # frame_idx ? [(track_id, box)]; {} when person
 cue OFF
 net: Tuple[str, float] # (axis, position): rally_gate auto-estimate or config
 frame_states: Dict[int, "ShuttleState"] # frame_idx ? state, {} until the shuttle-state
 detector ($3.5)
 base_stats: Dict[str, float] # motion stats already computed for the window

class WindowAnalyzer(ABC):

```

```

 name: str; version: str # registry key + column prefix; version ? provenance on
change
 feature_keys: List[str]; event_types: List[str] # declared columns/events (for the
harness)
 @abstractmethod
 def analyze(self, w: WindowSignals) -> AnalyzerResult: ...
...

```

The three roles an `AnalyzerResult` can fill:

```

| Role | Sink | Consumed by |
|---|---|---|
| *(a) features* (with confidence) | `candidate_segments.stats` as `<analyzer>__<key>` (+
`_confidence`) | the scoring model **and** the harness, as feature columns (§4) |
| *(b) events* | the `events` table (buffered on the candidate, flushed on confirmation,
§5-data) | history / highlights / audit |
| *(c) boundary hints* (optional) | stashed in `stats`; never auto-applied | the reconciler
(§6) / a future boundary-aware scorer |

```

### 3.2 Worked example ? the SERVICE-FAULT analyzer (window-scope)

The owner's held-shuttle bug has a sibling: a **service fault** where the serve hits the net and the shuttle comes to rest in the net region. Pure trajectory, GPU-free, no person cue needed:

```

- **Input:** the window's shuttle trajectory + the `rally_gate` net estimate.
- **Logic:** the shuttle **enters the net region** (`|coord ? net| < deadband`) **and** its
velocity
 decays to ? 0 there (comes to rest) ? high-probability net-hit service fault. Confidence
scales with
 how cleanly the halt sits inside the net band.
- **Emits:** a `service_fault` **signal** (value = presence, confidence = halt cleanliness) + a
`net_hit`
 event at the halt timestamp + a rally-**END** `BoundaryHint` at that timestamp.

```

Crucially: the analyzer does **\*\*not\*\*** drop the window. It **\*proposes\*** `service\_fault`; the scorer sees  
`service\_fault` + `service\_fault\_confidence` as two more columns and **\*\*learns\*\*** whether/how much to  
down-weight. The boundary hint is an audited proposal the reconciler may act on ? never a  
branch.

### 3.3 Injection point ? generalize `\_enrich\_candidate\_stats`

The FusionSegmenter hook (`fusion\_hybrid.py:70`) already turns "compute fixed features" into per-window enrichment. We generalize it to **\*\*"run each enabled analyzer"\*\*** ? additively, so the existing  
`net\_crossings`/person paths are unchanged:

```

```python
def _enrich_candidate_stats(self, cw, points, fps, frame_width, video_path, idx_cfg):
    stats = super()._enrich_candidate_stats(...)          # base motion keys
    stats["net_crossings"] = ...                          # Gate-A signal (kept)
    # ? person features when the person cue is on (kept) ?
    w = self._window_signals(cw, points, fps, frame_width, video_path, idx_cfg)
    for analyzer in analyzer_registry.enabled(idx_cfg):   # NONE enabled ? byte-identical to
control
        res = analyzer.analyze(w)
        for key, sig in res.signals.items():
            stats[f"{analyzer.name}__{key}"] = sig.value
            stats[f"{analyzer.name}__{key}_confidence"] = sig.confidence
        stats.setdefault("_events", []).extend(e.__dict__ for e in res.events)          # flushed
in Phase 4
        stats.setdefault("_boundary_hints", []).extend(h.__dict__ for h in res.boundary_hints)
    return stats

```

...

Flag-gated, ****default OFF**** via the existing ``indexing.fusion.cue_flags`` (a sibling ``analyzer_flags`` map keyed by ``analyzer.name``). With every analyzer OFF, the loop is empty ? candidate rows are byte-identical to the substrate ? the no-regression invariant holds (§7). ``net_crossings`` and the person features are the ****prototype analyzers**** the seam retrofits first (proving it with zero behavior change); every new cue after that is a pure add.

3.4 Session-scope analyzers ? the WARM-UP / PRACTICE example

Some structure is invisible at the window scope. ****warm-up**** looks exactly like a rally in any single window ? shuttle in flight, two players, net crossings, sustained motion ? which is **why** it is a notorious false-positive source. It is only distinguishable by ****timeline structure****: a long contiguous active stretch **without** the rally?dead-time?rally cadence of scored play, usually (not always) at the start.

So warm-up needs a producer that sees the ****whole sequence of windows****, not one window. That is a sibling tier ? a ``SessionAnalyzer`` that runs once per video (a second pass, after per-window enrichment) and emits a labeled ****span**** plus a per-window ****membership feature****:

```
```python
@dataclass(frozen=True)
class Span:
 kind: str # "warmup" | "practice" | "break" | "game" | ?
 start: float; end: float; confidence: float; reason: str

@dataclass(frozen=True)
class SessionResult:
 spans: List[Span] = field(default_factory=list)
 # window_index ? signals merged back into THAT window's stats (e.g. {"in_warmup":
 Signal(0.8, 0.6)})
 window_features: Dict[int, Dict[str, Signal]] = field(default_factory=dict)

class SessionAnalyzer(ABC):
 name: str; version: str
 span_kinds: List[str]; feature_keys: List[str]
 @abstractmethod
 def analyze(self, windows: List[WindowSignals], session: "SessionContext") -> SessionResult:
 ...
 ...
```

```
`WarmupAnalyzer` (worked):
- **v0 (free, no new model):** find the **leading contiguous high-activity span** with low
inter-window
 dead-time and no serve-reset cadence, over the persisted window timings we already have. Emit
 `Span(kind="warmup", ?)` + per-window `in_warmup` = overlap fraction of the window with the
 span.
- **v1 (learned):** fed by the **shuttle-state cadence** (§3.5) / the **audio thwack** cadence ?
warm-up =
 continuous hits *without* the point-ending pause rhythm; generalizes to mid-session practice /
between-
 game knock-ups (drop the position prior).
```

**\*\*You framed the consumption exactly right: "so the scorer may consume it."\*\*** ``in_warmup`` is a **\*\*feature column\*\***; the scorer *\*learns\** to down-weight windows inside the span. **\*\*Warm-up** is the poster child for

no-special-casing\*\*: the naive version is `if t < warmup\_end: drop` ? a hand-coded branch hanging off a fuzzy boundary; a 20-second error \*\*deletes a real opening rally\*\*. As a weighted feature, a wrong boundary just slightly mis-weights a couple of windows the other cues rescue. The span is \*also\* persisted as evidence a reconciler (§6) may use to down-weight a whole cluster, but the primary path is feature ? scorer.

\*\*It is measurable today with no new labels:\*\* the golden labels already treat warm-up knock-ups / pre-serve as \*\*dead time = negatives\*\* (Roora starts a rally at serve-contact; fusion §1 names "warm-up knock-ups" as a precision target). So an `in\_warmup` feature is directly A/B-able on the 6 golden videos for \*\*precision\*\* lift ? that is the promotion gate.

### 3.5 Consuming a richer detector ? SHUTTLE-STATE (frame-scope producer)

A natural new \*\*detector\*\*: a per-frame shuttle-state classifier ? `in\_air` / `racket\_contact` / `inactive`. Note the scope: it \*\*processes frames ? a new per-frame signal\*\*, like WASB, so it is a

\*\*frame-scope producer in the detection layer\*\* ? its spec belongs in [MULTI\_SIGNAL\_FUSION\_PLAN.md](MULTI\_SIGNAL\_FUSION\_PLAN.md) (it is a new `RallyCue`), \*\*not\*\* a `WindowAnalyzer`.

This doc records \*\*how the framework consumes it\*\* once it exists:

- a \*\*`WindowAnalyzer`\*\* turns the persisted per-frame state into window features ?  
`racket\_contact\_count`  
(? shot cadence), `inactive\_ratio` (? dead-time) ? that flow into the scorer as columns;
- a \*\*`SessionAnalyzer`\*\* (warm-up, §3.4) reads its cadence;
- the \*\*reconcilers\*\* (§6) get \*stronger\* evidence: `trailing\_dead\_time` / `over\_merged`  
currently infer  
inactivity from a velocity threshold; a learned `inactive` state is a more direct, occlusion-robust  
evidence source for "truncate / split here."

\*\*Stage it (don't build the model first).\*\* Most of the state is derivable for free from signals we already

have ? `in\_air` vs `inactive` ? WASB velocity + visibility; `racket\_contact` ? trajectory  
\*\*direction

reversals\*\* (fusion §4: "each reversal ? a shot") \*\*and\*\* the scaffolded \*\*audio thwack\*\*  
(`backend/pipeline/audio/hit\_detector.py`, ADR-007's #1 cue). So \*\*v0 = a derived  
`shuttle\_state` analyzer\*\*

over existing signals (no new model); measure it. \*\*The dedicated per-frame vision model earns its keep

only where v0 hits a ceiling: WASB trajectory gaps\*\* ? occlusion, motion blur, fast smashes ?  
where there

is no `(x,y)` to threshold but an appearance model could still call `racket\_contact`. \*\*Label  
blocker:\*\* we

have \*\*zero\*\* per-frame state labels (golden labels are rally intervals); a dedicated model  
needs weak/

distilled labels or a new (minors-gated) annotation task ? the same gate as feature-level  
fusion. If you can

weak-label state from existing signals, prefer the derived analyzer until the gap-filling case  
is proven.

---

## 4. How signals reach the scorer ? columns, not branches

The persisted `stats` columns are exactly what the harness already lifts.

`load\_video\_candidates(...,  
feature\_names=[...])` (`experiment.py:429`) reads each name out of the row's `stats` JSON into a

```
feature
column; `predict()` (`experiment.py:203`) gates/scores on those columns and merges. So a new
producer's
signal becomes a scorer input by **adding its column name to a `Variant`'s `feature_names`** ?
no new
scoring code, no branch:
```

```
`python
```

## A learned Variant that lets the scorer weight the new signals (service-fault + warm-up)

```
fusion_plus = Variant(name="fusion+analyzers",
 feature_names=[*SHUTTLE_FEATURES, "person__players_in_court",
 "service_fault__service_fault",
 "service_fault__service_fault_confidence",
 "warmup__in_warmup", "warmup__in_warmup_confidence"])
compare(golden_videos, CONTROL, fusion_plus) # LOVO-honest B ? A on the 6 golden videos
`python`
```

The model \*\*learns the weights\*\* ? down-weight a window with a high-confidence `service\_fault` or high `in\_warmup`; up-weight ?2 in-court players + multiple net-crossings. \*\*What the scorer cannot see, it cannot special-case; what it can see, it weights from data.\*\* Same logistic regression as ADR-004; we add columns, never a net.

```

```

## 5. Persistence

- \*\*Features ? `candidate\_segments.stats`\*\* (`database.py:240`). Namespaced `>\_\_<key>` (+ `\_confidence`). Persisted \*\*as a shadow\*\* even when the served decision ignores them, so any future `Variant` re-scores offline (EXPERIMENT\_HARNESS §5.2). `version`/`name` recorded for provenance.
- \*\*Spans ? window membership + a session record.\*\* A `SessionAnalyzer`'s per-window `in\_<span>` features ride in each window's `stats` (so they re-score like any feature); the span list itself (start/end/kind/confidence) is a small per-video record (a JSON sidecar first; a table only if it grows).
- \*\*Events ? the `events` table\*\* (`database.py:227`, `add\_event(segment\_id, ts, type, player, details)`). Its FK is to `play\_segments`, which exist only \*\*after\*\* the AI handoff confirms a candidate. So analyzer events are \*\*buffered on the candidate\*\* (`stats["\_events"]`) and \*\*flushed\*\* to `events` in Phase 4, when `\_candidate\_id` ? `segment\_id` is linked (`trajectory\_hybrid.py:283`). Candidate-level discoveries survive for offline re-scoring; the `events` table stays "events on confirmed segments."
- \*\*Corrections record\*\* (§6) ? report-first means corrections are \*\*recorded, never silent.\*\* Phase one: a JSON sidecar shaped like `experiment\_leaderboard.json` (one row per issue: segment, kind, evidence, before?after, reason, timestamp). When auto-apply is promoted, add a `corrections` table via the `PRAGMA user\_version` migration ladder (`database.py:88`, `if version < 3:`) so applied corrections are queryable and reversible.

```

```

## 6. The backward seam ? the reconciliation pass (a linter for rally decisions)

Given **\*\*predicted segments + the persisted per-frame/per-window/per-session signals\*\***, detect where the **\*\*DECISION contradicts the EVIDENCE\*\*** and propose corrections. This is the backward dual of the analyzers:  
 analyzers add evidence **\*before\*** the decision; reconcilers audit the decision **\*against\*** evidence **\*after\*** it.

## 6.1 Contracts (proposed `backend/pipeline/reconcilers/`)

```
```python
@dataclass(frozen=True)
class Issue:
    segment_index: int; kind: str          # kind = rule id, e.g. "trailing_dead_time"
    evidence: Dict[str, Any]               # the signal values that triggered it (auditable)
    before: Interval; after: Optional[List[Interval]] # None=drop · [iv]=corrected ·
[iv1,iv2]=split
    reason: str; confidence: float

@dataclass(frozen=True)
class Correction:
    issues: List[Issue]; corrected: List[Interval] # the FULL corrected set (idempotent)

class ConsistencyRule(ABC):
    name: str; precedence: int             # FIXED ordering; lower runs first
    @abstractmethod
    def check(self, seg: Interval, signals: "SegmentSignals") -> Optional[Issue]: ...

def reconcile(segments, signals, rules) -> Correction: ... # fixed precedence; idempotent
```
```

## 6.2 Worked rules (fixed precedence)

| # | Rule                 | Evidence contradiction                                                                               | Correction                                                                      |
|---|----------------------|------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------|
| 1 | `trailing_dead_time` | segment still "in rally" at its end, but the shuttle went inactive earlier                           | <b>**truncate**</b> to the last-active frame <b>*(the owner's misfire fix)*</b> |
| 2 | `no_net_crossing`    | `net_crossings < 1` ? nothing ever crossed the net                                                   | <b>**drop**</b> (or flag)                                                       |
| 3 | `over_merged`        | a dead-time gap <b>**&gt;**</b> `merge_gap` of shuttle-inactivity sits <b>**inside**</b> one segment | <b>**split**</b> at the gap                                                     |
| 4 | `boundary_slack`     | first sustained shuttle activity / first net-crossing is well inside the start                       | <b>**tighten**</b> start to serve-contact                                       |
| 5 | `inside_warmup`      | the segment lies inside a high-confidence `warmup`/`practice` span (\$3.4)                           | <b>**flag**</b> (drop only under apply-flag)                                    |

Each rule reads only **\*\*persisted evidence\*\*** (trajectory activity per frame, `net\_crossings`, person occupancy, analyzer signals/events, session spans, shuttle-state) ? it is a deterministic linter, not a re-detector. As richer producers land (shuttle-state, warm-up), rules get **\*\*stronger evidence\*\***, not new branches in the decision path.

## 6.3 Report-first, apply-optional (non-negotiable)

`reconcile()` returns **\*\*both\*\*** an issues report **\*\*and\*\*** a corrected set. It **\*\*NEVER silently rewrites\*\*** (auditability + the repo's honesty/attribution working agreement). The pipeline either:

- **\*\*default\*\*** persists the report and **\*\*surfaces it for human review\*\*** ? production output unchanged; or
- **\*\*flagged, `reconcile.apply`\*\*** applies the corrected set, **\*\*recording every change\*\*** in the corrections record (before?after, rule, evidence). Rollback = flip the flag.

**\*\*It is a `Variant` transformation, so its lift is A/B-measurable BEFORE auto-apply.\*\*** A reconciler is a pure post-`predict` stage: `reconcile(predict(...), signals, rules).corrected`. Wrapping that as a `RECONCILED` variant lets `compare(CONTROL, RECONCILED)` report the LOVO lift on the 6 golden videos ([QUALITY\_ITERATIONS.md](QUALITY\_ITERATIONS.md) corpus) ? promote to auto-apply **\*\*only\*\*** on measured lift.

**\*\*Idempotent; fixed rule precedence.\*\*** `reconcile(reconcile(x)) == reconcile(x)` (truncating an already-truncated segment is a no-op); precedence is a fixed table so the result is deterministic and order-independent of input segment order.

---

## 7. No-regression discipline (identical to the harness contract)

The same four guarantees EXPERIMENT\_HARNESS.md §6 makes ? applied to every producer and the reconciler:

- **\*\*Default OFF.\*\*** Analyzers gate on `analyzer\_flags` (default empty); the reconciler is report-first (apply gate default off). Merged code **\*\*cannot\*\*** change served behavior until explicitly enabled.
- **\*\*Shadow-persist, then promote.\*\*** Signals/spans/issues are persisted even while ignored, so a `Variant` re-scores them offline; a signal becomes a *\*decision input\** only when a Variant that references it wins LOVO. (EXPERIMENT\_HARNESS §5.2's shadow mode.)
- **\*\*Promotion only via the eval gate.\*\*** `run\_regression.py` against the golden `baseline.json` (exit 1 blocks); the experiment leaderboard records the spec hash + per-video + LOVO deltas + label-set hash.
- **\*\*Rollback = flip a config flag, never `git revert`.\*\*** The pinned `CONTROL` variant is always present.

The keystone test (extends the harness's existing invariant): **\*\*all analyzer flags OFF + reconciler report-only ? candidate rows AND predictions byte-identical to control.\*\***

---

## 8. How it composes with fusion P2

- **\*\*Window-analyzers plug into fusion P2's enrichment hook.\*\*** `FusionSegmenter.\_enrich\_candidate\_stats` (`fusion\_hybrid.py:70`) is the exact injection point §3.3 generalizes; `net\_crossings` and the person features become the first registered analyzers (no behavior change). New cues are add-a-producer, per ADR-007 / MULTI\_SIGNAL\_FUSION\_PLAN §2.
- **\*\*Session-analyzers run as a second pass\*\*** over the per-window stats the same hook persists.
- **\*\*The shuttle-state detector is a new fusion `RallyCue`\*\*** (frame-scope producer) ? it belongs in the fusion plan; this framework only *\*consumes\** it (§3.5).
- **\*\*Reconcilers consume fusion P2's persisted signals\*\*** ? the same `candidate\_segments.stats` plus the persisted trajectory; no new detection.
- Everything inherits fusion P2's control-reproducing defaults, so the fusion segmenter with all flags off still seeds and hands off identically to `wasb\_hybrid`.

---



## 9. Trade-offs & risks

- **Deterministic vs learned signals.** A producer may emit a verdict-shaped signal (`service_fault=1.0``), but we persist it as a **feature**, so the scorer still arbitrates. Risk: a confidently-wrong cue.  
Mitigation: the **confidence column** + the **LOVO gate** ? a bad signal earns ? 0 weight, never promotes.
- **The temporal confound is the hard core.** Warm-up-vs-play, multi-court, and held-shuttle all look identical in any single frame / frequency band ? `QUALITY_ITERATIONS`' standing lesson: "the dead-time confusion is rally-vs-non-rally activity in the same place and frequency band... only the learned model's temporal features crack it." This is *why* warm-up is a **feature the model weights**, not a heuristic span-cutter, and why session-scope structure earns its own tier.
- **Scope discipline.** Put a producer at its *natural* scope: shuttle-state is frame-scope (a detector), not a window analyzer; warm-up is session-scope, not per-window. Mis-scoping forces a producer to either recompute detection or miss the structure it needs.
- **Ordering / idempotence (reconcilers).** Fixed precedence table + an idempotence test (`reconcile?reconcile == reconcile``).
- **Signal, not verdict.** A producer that *decides* (returns keep/drop) re-introduces special-casing; reviews must reject "analyzer returns a boolean keep/drop."
- **Report-first auditability.** Reconcilers never silently mutate output; every applied change is recorded.
- **Small-N feature discipline.** Few, scale/translation-invariant columns (counts/ratios/presence, never raw coords ? `MULTI_SIGNAL_FUSION_PLAN §3.2`); the model stays a logistic regression, not a net.
- **Labels gate the learned producers.** Shuttle-state and feature-level fusion need labels we don't yet have; warm-up is measurable *now* because golden treats it as negatives. Know which producers are gated.
- **Cost & licensing.** Pure analyzers are GPU-free; person- and shuttle-state-dependent producers pull torch / a second WSL pass ? gate behind their cue flags; all deps stay MIT/Apache/BSD (`ADR-002/005`).
- **Event-table timing.** The buffer-then-flush wiring (§5) is the one non-obvious integration detail.

---

## 10. Generality & research positioning

This is written as a **framework**, not a badminton patch. Read sport-agnostically it is: *a registry-based, multi-scope **signal-fusion** architecture for **weakly-supervised temporal event/segment detection**, with a learned arbiter and an **auditable post-hoc correction** layer.* The pieces that generalize:

- **Multi-scope producers** ? one feature table ? a learned arbiter. The frame/window/session taxonomy and the "signal-not-verdict" rule are domain-independent. Any temporal problem with **weak, multi-modal cues at different time scales** (broadcast sports event spotting, surgical-phase recognition, call-center / meeting segmentation, activity logs) fits the same shape: add a producer, add a column, let the arbiter weight it, gate promotion on held-out lift.
- **Detection ? decision** (persist once, re-score forever). Decouples the expensive, hard-to-reproduce

perception from the cheap, deterministic decision ? the property that makes the whole thing experimentally honest and cheap to iterate.

- **Auditable reconciliation as a first-class layer.** Most pipelines bury post-hoc fixes as silent if-branches; treating consistency-vs-evidence as a *report-first, A/B-measured* pass is the transferable methodological idea (and the human-in-the-loop / safety story).
- **The honest measuring stick.** Interval-level, scarce ground truth + *leave-one-VIDEO-out* (group by match, never a held-out window from the same video) is the discipline any small-corpus temporal study needs to avoid flattering itself.

**What is defensibly contributed vs. what is standard.** Multi-modal *late fusion*, weak supervision, and learning a linear/logistic arbiter over hand-designed features are all standard. The *combination* offered here ? (a) a uniform producer-registry across *three temporal scopes*, (b) a hard *no-branch* decision invariant that makes extension monotone, (c) a *report-first, harness-measured reconciliation* layer, and (d) the *detection?decision persistence* substrate that makes every experiment a zero-risk offline re-score ? is what would be written up, with rally detection as the motivating hard problem.

**Honesty caveat (attribution working agreement).** No external results or citations are asserted here. A real write-up still owes: a proper *related-work survey* (temporal action segmentation/detection, sports event spotting, multimodal fusion, weak supervision ? to confirm which claims are actually novel), *strong baselines*, *per-producer ablations*, and reporting on *>1 sport / camera* to substantiate the generality claim. Today the corpus is *6 matches* (QUALITY\_ITERATIONS) ? enough for direction-checks and LOVO sign-tests, *not* enough to claim a general model; corpus growth is the highest-ROI unblock. Treat §10 as a *positioning sketch*, not a results section.

---

## 11. Phased execution (owner-gated; ONE PR per slice, off `master`, default OFF)

- **A0 ? this doc**, indexed in `docs/README.md`, cross-linked from the peers + ADR-007. *(this deliverable; docs-only.)*
- **A1 ? analyzer contracts + registry**  
 (`WindowAnalyzer` / `AnalyzerResult` / `Signal` / `AnalyzerEvent` / `BoundaryHint` + `analyzer\_registry`); generalize `\_enrich\_candidate\_stats` to iterate enabled analyzers (none enabled ? byte-identical). Retrofit `net\_crossings` + person features as the first two analyzers.
- **A2 ? service-fault analyzer** behind `analyzer\_flags.service\_fault` (default OFF); signal + `net\_hit` event + end boundary hint; shadow-persisted; offline `compare(CONTROL, +service\_fault)` LOVO.
- **A3 ? event flush wiring** (buffer analyzer events on candidates ? `add\_event` on confirmation).
- **A4 ? `SessionAnalyzer` tier + `WarmupAnalyzer` v0** (leading-span heuristic over window timings), default OFF; persists `warmup\_\_in\_warmup` + the span record; measured for *precision* lift on golden.
- **A5 ? derived `shuttle\_state` analyzer v0** (velocity + visibility + direction-reversals (+ audio); *no new model*); emits `racket\_contact\_count` / `inactive\_ratio`; measured LOVO. *(The dedicated*

```

 per-frame shuttle-state DETECTOR is a separate MULTI_SIGNAL_FUSION_PLAN `RallyCue` item ?
gated on
 per-frame labels; build only if v0 is blocked by trajectory gaps.)*
- **R1 ? reconciler contracts + registry + `reconcile()`**
(`ConsistencyRule`/`Issue`/`Correction`) with
 the worked rules; pure, idempotent, unit-tested; **not** wired to apply (report-only).
- **R2 ? reconciler as a `Variant` transform**: optional post-`predict` stage; `compare(CONTROL,
 RECONCILED)` LOVO; report-only in production.
- **R3 ? corrections persistence + opt-in apply**: JSON sidecar first, then the `corrections`
table
 (`user_version < 3`); apply behind `reconcile.apply`; **rollback = flip the flag.**

```

Each PR is gated by the no-regression invariant (\$?) + `run\_regression.py`. Order is a suggestion; A4/A5/R1 are independent and can be reprioritized on owner steer.

---

## 12. Verification

```

- **Offline, runnable in this dev container (no GPU/torch/cv2):** pure analyzer math on
synthetic
 trajectories + person boxes (service-fault fires on a synthetic net-halt, silent on a clean
rally;
 `WarmupAnalyzer` v0 finds a synthetic leading knock-up span; derived `shuttle_state` matches a
scripted
 velocity/visibility sequence); the **no-regression invariant** (all flags OFF + reconciler
report-only ?
 candidate rows + predictions byte-identical to control); the **reconciler idempotence** test
 (`reconcile?reconcile == reconcile`) and fixed-precedence determinism; `compare(CONTROL,
variant)`
 deterministic on synthetic candidate rows.
- **Machine-side (owner, GPU/WSL ? full pipeline can't run here):** persist analyzer signals on
a real
 clip; `compare(CONTROL, +service_fault)`, `compare(CONTROL, +warmup)`, `compare(CONTROL,
RECONCILED)`
 LOVO on the 6 golden videos; confirm the owner's **trailing-dead-time misfire** truncates, the
held-shuttle / net-hit service feed is down-weighted, and **warm-up windows** lose
precision-mass;
 spot-check `net_hit` events on confirmed segments.
- **Docs (A0):** links resolve; indexed in `docs/README.md`; cross-linked from
 MULTI_SIGNAL_FUSION_PLAN.md,
EXPERIMENT_HARNESS.md,
HOW_RALLY_DETECTION_WORKS.md, and ADR-007.

```

### 94. user (2026-06-13T19:44:48.727Z)

```

1 # Analyzers & reconcilers ? two extensibility seams for rally-detection quality
2
3 **Status:** **DESIGN ? owner-gated, NOT started.** A peer to

```

### 95. assistant (2026-06-13T19:44:47.439Z)

The file changed on disk (autocrlf likely re-touched it during the git ops). Let me re-read it, then write.

### 96. assistant (2026-06-13T19:47:05.781Z)

## A signal-fusion framework for rally detection ? analyzers, scopes & reconcilers

```

Status: **DESIGN ? owner-gated, NOT started.** A peer to
EXPERIMENT_HARNESS.md and
MULTI_SIGNAL_FUSION_PLAN.md.

```

This doc specs a small **framework** for evolving rally-detection quality: producers that emit **signals** at three temporal **scopes** (frame / window / session), a **scorer** that *learns to weight* them, and an auditable **reconciliation** pass that lints the decision against the evidence. The goal is to keep piling on clever cues **without the pipeline becoming a pile of special cases.** Docs-only; no implementation code. Each phase below is a separate, owner-gated PR.

```

> Companion docs: the detection-layer producers (shuttle / person / audio / shuttle-state) live in
> MULTI_SIGNAL_FUSION_PLAN.md (`RallyCue` = swappable producer; the
> `_enrich_candidate_stats` hook; scale/translation-invariant feature discipline, §3.2); the no-regression
> A/B + LOVO measurement these ride is EXPERIMENT_HARNESS.md
> (`Variant`, `compare`,
> `compare_unlabeled`); the scorer is the distilled logistic regression (ADR-004,
> `backend/eval/classifier.py`); the modular rally layer is ADR-007
> (archives/decisions/DECISIONS.md); the 2-layer mental model is
> HOW_RALLY_DETECTION_WORKS.md; the live quality log is
> QUALITY_ITERATIONS.md.

> Framing. Rally detection is a genuinely hard temporal problem ? weak, multi-modal cues; interval-level
> ground truth that is scarce and expensive; a confound (warm-up / knock-ups / multi-court / held-shuttle)
> that looks identical to real play in any single frame or frequency band. This doc is written as a
> reusable framework, not a badminton-specific patch: the contracts are sport-agnostic and the design
> is meant to generalize to other weakly-supervised temporal-segmentation problems (§10). Where that
> ambition touches research claims, §10 says plainly what is and isn't yet substantiated.

```

---

## 1. The spine ? NO special-casing in the decision path

This is the one principle everything else serves:

```

> Every producer (analyzer, detector, reconciler) emits a SIGNAL carrying a value + a confidence/presence.
> The SCORING MODEL consumes those signals and learns to weight them. We do NOT add hand-coded `if/else`
> branches into the decision.

```

**Producers propose; the scorer arbitrates; the experiment harness + golden set decide what gets promoted.**

That separation is what lets us add a net-hit service fault, a let-cord, a double-bounce, a shuttle-out cue, a *warm-up span*, a *shuttle-state* ? each as `"register a producer + a feature column,"` never `"add a branch."` A signal that doesn't earn its keep stays at weight ? 0; it costs nothing because it is an unused column, not a live ``if``.

The pattern already exists in miniature: ``experiment.predict()`` (``backend/eval/experiment.py:203``) treats ``net_crossings`` and ``players_in_court`` as **feature columns** ? ``_gate_mask`` is the trivial scorer (threshold a column), ``LogisticRegression`` (ADR-004) the richer learned scorer (weight all

```

columns), a
future scorer richer still. The decision path is *signals ? scorer*, with **zero** `if
service_fault:
drop`. Analyzers feed it more columns; reconcilers run *outside* it as an audited post-pass.
...
DETECTION ? frame-scope producers (expensive, GPU/WSL, run ONCE)
 WASB shuttle (x,y,v,vis) · person boxes+tracks · audio hits · shuttle-
state(in_air/contact/inactive)*
 ? per-frame-aligned signals (+ net estimate, window bounds)
 ?
 shuttle-seeded candidate windows (FusionSegmenter, P2)
 ?
ANALYZERS ? decision-layer producers (pure, offline, re-scorable) [forward seam]
 ?? window-scope (WindowAnalyzer): one candidate ? features + events + boundary hints e.g.
service-fault
 ?? session-scope (SessionAnalyzer): whole timeline ? labeled spans + per-window membership
e.g. warm-up
 ? every signal ? candidate_segments.stats as feature columns
 ?
SCORING / DECISION ? signals ? learned scorer (weights, NO branches)
 _gate_mask (trivial) + LogisticRegression (ADR-004) + harness-pinned model
 ? predicted rally segments
 ?
RECONCILERS ? backward seam, "linter for rally decisions" (report-first, idempotent)
 ConsistencyRule registry: decision-vs-evidence conflicts ? Issue + corrected set
 ?
 stats · events · spans · corrections record · highlights EXPERIMENT HARNESS gates EVERY
promotion
 (* shuttle-state = a richer frame detector; its PRODUCER lives in MULTI_SIGNAL_FUSION_PLAN,
consumed here ? §3.5)
...

```

## 2. The framework ? signals, scopes, and who arbitrates

The architecture has exactly three kinds of moving part. Capability is added by registering a producer; the decision logic never grows a branch.

**\*\*Producers emit signals at three temporal SCOPES.\*\*** A signal at a coarser scope is computed by aggregating finer-scope signals; everything ultimately lands as a per-candidate feature column (or an event/span the scorer or a human consumes).

| Scope       | Producer                 | Sees                                         | Emits                                          | Cost          | Examples                                                                     |
|-------------|--------------------------|----------------------------------------------|------------------------------------------------|---------------|------------------------------------------------------------------------------|
| ---         | ---                      | ---                                          | ---                                            | ---           | ---                                                                          |
| **frame**   | detector (`RallyCue`)    | raw frames / audio                           | per-frame signal                               | **expensive** | (GPU/WSL)   WASB shuttle, person boxes, audio hits, **shuttle-state** (§3.5) |
| **window**  | `WindowAnalyzer` (§3.1)  | one candidate window's aligned frame-signals | features (+conf), events, boundary hints       | cheap, pure   | service-fault, double-bounce, let-cord, shuttle-out                          |
| **session** | `SessionAnalyzer` (§3.4) | the whole timeline of windows                | labeled spans + per-window membership features | cheap, pure   | **warm-up / practice**, game/break phases                                    |

**\*\*The scorer arbitrates.\*\*** It is the only place a decision is made, and it makes it by **\*weighting columns\***, never branching. Today: `\_gate\_mask` (threshold) + `LogisticRegression` (learned weights, ADR-004). The confidence column on every signal lets the scorer discount an uncertain producer instead of trusting a binary verdict ? the small-N-safe way to fold in a deterministic-ish cue.

**\*\*Reconcilers audit.\*\*** After the scorer decides, the backward seam (§6) lints the decision against the persisted evidence and proposes corrections ? report-first, never silent.

**\*\*Four invariants make this safe to grow\*\*** (the load-bearing design claims):

1. **\*\*Signal, not verdict\*\*** ? every producer's output is a weighted feature; the decision path has no producer-specific branch. ? unbounded cues without combinatorial special-casing.
2. **\*\*Expensive detection ? cheap decision\*\*** ? frame-scope detection runs once and is persisted; window/session signals and the scorer are pure re-scoring (EXPERIMENT\_HARNESS §2). ? experiments are offline, deterministic, zero production risk.
3. **\*\*Report-first reconciliation\*\*** ? corrections are recorded and reviewable, never applied silently. ? human-in-the-loop, auditable.
4. **\*\*Promote only on measured held-out lift\*\*** ? LOVO on the golden set + the `run\_regression.py` gate; default OFF; rollback = config flip. ? no-regression evolution.

---

### 3. The forward seam ? injectable analyzers

Pluggable producers ? registered like `segmenter\_registry` (`backend/pipeline/segmenters/base.py:35`) ? that run over the **\*\*per-frame-aligned signals already produced\*\***: the WASB shuttle trajectory (`frame/.visible/.x/.y` + velocity), the person boxes per frame (`PersonDetector.detections\_per\_frame`, `person\_detector.py:256`), the net line (`rally\_gate.estimate\_play\_axis`, camera-agnostic), and the window bounds. Analyzers are **\*\*pure\*\*** ? no I/O, no torch/cv2 ? mirroring `fusion\_features.py`'s discipline, so their math is unit-testable in the GPU-free dev container.

#### 3.1 Contracts (JSON-serializable; proposed `backend/pipeline/analyzers/`)

```
```python
@dataclass(frozen=True)
class Signal:
    # one scalar cue + how sure the producer is ? never a verdict
    value: float; confidence: float = 1.0    # persisted as `<producer>__<key>` (+
`_confidence`), [0,1]

@dataclass(frozen=True)
class AnalyzerEvent:
    # ? the `events` table on confirmation (§5)
    type: str; timestamp: float; confidence: float = 1.0    # "net_hit"/"double_bounce"/? ; src-absolute s
    player: Optional[str] = None; details: Dict[str, Any] = field(default_factory=dict)

@dataclass(frozen=True)
class BoundaryHint:
    # OPTIONAL deterministic proposal ? never auto-applied
    kind: Literal["start", "end"]; timestamp: float; confidence: float; reason: str

@dataclass(frozen=True)
class AnalyzerResult:
    signals: Dict[str, Signal] = field(default_factory=dict)
    events: List[AnalyzerEvent] = field(default_factory=list)
    boundary_hints: List[BoundaryHint] = field(default_factory=list)

@dataclass
class WindowSignals:
    # read-only pre-produced bundle (no I/O inside the analyzer)
    start: float; end: float; fps: float; frame_width: float; frame_height: float
```

```

    shuttle: List["TrajectoryPoint"]      # in-window points (.frame/.visible/.x/.y)
    persons: Dict[int, List[Tuple[int, BBox]]] # frame_idx ? [(track_id, box)]; {} when person
cue OFF
    net: Tuple[str, float]                # (axis, position): rally_gate auto-estimate or config
    frame_states: Dict[int, "ShuttleState"] # frame_idx ? state; {} until the shuttle-state
detector (§3.5)
    base_stats: Dict[str, float]          # motion stats already computed for the window

class WindowAnalyzer(ABC):
    name: str; version: str                # registry key + column prefix; version ? provenance on
change
    feature_keys: List[str]; event_types: List[str] # declared columns/events (for the
harness)
    @abstractmethod
    def analyze(self, w: WindowSignals) -> AnalyzerResult: ...
...

```

The three roles an `AnalyzerResult` can fill:

```

| Role | Sink | Consumed by |
|---|---|---|
| *(a) features* (with confidence) | `candidate_segments.stats` as `<analyzer>__<key>` (+
`_confidence`) | the scoring model and the harness, as feature columns (§4) |
| *(b) events* | the `events` table (buffered on the candidate, flushed on confirmation, §5) |
history / highlights / audit |
| *(c) boundary hints* (optional) | stashed in `stats`; never auto-applied | the reconciler
(§6) / a future boundary-aware scorer |

```

3.2 Worked example ? the SERVICE-FAULT analyzer (window-scope)

The owner's held-shuttle bug has a sibling: a **service fault** where the serve hits the net and the shuttle comes to rest in the net region. Pure trajectory, GPU-free, no person cue needed:

```

- Input: the window's shuttle trajectory + the `rally_gate` net estimate.
- Logic: the shuttle enters the net region (`|coord ? net| < deadband`) and its
velocity
  decays to ? 0 there (comes to rest) ? high-probability net-hit service fault. Confidence
scales with
  how cleanly the halt sits inside the net band.
- Emits: a `service_fault` signal (value = presence, confidence = halt cleanliness) + a
`net_hit`
  event at the halt timestamp + a rally-END `BoundaryHint` at that timestamp.

```

Crucially: the analyzer does **not** drop the window. It **proposes** `service\_fault`; the scorer sees `service\_fault` + `service\_fault\_confidence` as two more columns and **learns** whether/how much to down-weight. The boundary hint is an audited proposal the reconciler may act on ? never a branch.

3.3 Injection point ? generalize `\_enrich\_candidate\_stats`

The FusionSegmenter hook (`fusion\_hybrid.py:70`) already turns "compute fixed features" into per-window enrichment. We generalize it to **"run each enabled analyzer"** ? additively, so the existing `net\_crossings`/person paths are unchanged:

```

```python
def _enrich_candidate_stats(self, cw, points, fps, frame_width, video_path, idx_cfg):
 stats = super()._enrich_candidate_stats(...) # base motion keys
 stats["net_crossings"] = ... # Gate-A signal (kept)
 # ? person features when the person cue is on (kept) ?
 w = self._window_signals(cw, points, fps, frame_width, video_path, idx_cfg)
 for analyzer in analyzer_registry.enabled(idx_cfg): # NONE enabled ? byte-identical to

```

```

control
 res = analyzer.analyze(w)
 for key, sig in res.signals.items():
 stats[f"{analyzer.name}__{key}"] = sig.value
 stats[f"{analyzer.name}__{key}_confidence"] = sig.confidence
 stats.setdefault("_events", []).extend(e.__dict__ for e in res.events) # flushed
in Phase 4
 stats.setdefault("_boundary_hints", []).extend(h.__dict__ for h in res.boundary_hints)
return stats
...

```

Flag-gated, **\*\*default OFF\*\*** via the existing ``indexing.fusion.cue_flags`` (a sibling ``analyzer_flags`` map keyed by ``analyzer.name``). With every analyzer OFF, the loop is empty ? candidate rows are byte-identical to the substrate ? the no-regression invariant holds (§7). ``net_crossings`` and the person features are the **\*\*prototype analyzers\*\*** the seam retrofits first (proving it with zero behavior change); every new cue after that is a pure add.

### 3.4 Session-scope analyzers ? the WARM-UP / PRACTICE example

Some structure is invisible at the window scope. **\*\*Warm-up\*\*** looks exactly like a rally in any single window ? shuttle in flight, two players, net crossings, sustained motion ? which is *\*why\** it is a notorious false-positive source. It is only distinguishable by **\*\*timeline structure\*\***: a long contiguous active stretch *\*without\** the rally?dead-time?rally cadence of scored play, usually (not always) at the start.

So warm-up needs a producer that sees the **\*\*whole sequence of windows\*\***, not one window. That is a sibling tier ? a ``SessionAnalyzer`` that runs once per video (a second pass, after per-window enrichment) and emits a labeled **\*\*span\*\*** plus a per-window **\*\*membership feature\*\***:

```

```python
@dataclass(frozen=True)
class Span:
    kind: str                # "warmup" | "practice" | "break" | "game" | ?
    start: float; end: float; confidence: float; reason: str

@dataclass(frozen=True)
class SessionResult:
    spans: List[Span] = field(default_factory=list)
    # window_index ? signals merged back into THAT window's stats (e.g. {"in_warmup":
Signal(0.8, 0.6)})
    window_features: Dict[int, Dict[str, Signal]] = field(default_factory=dict)

class SessionAnalyzer(ABC):
    name: str; version: str
    span_kinds: List[str]; feature_keys: List[str]
    @abstractmethod
    def analyze(self, windows: List[WindowSignals], session: "SessionContext") -> SessionResult:
...

```

****`WarmupAnalyzer`**** (worked):

- ****v0 (free, no new model):**** find the ****leading contiguous high-activity span**** with low inter-window dead-time and no serve-reset cadence, over the persisted window timings we already have. Emit ``Span(kind="warmup", ?)`` + per-window ``in_warmup`` = overlap fraction of the window with the span.
- ****v1 (learned):**** fed by the ****shuttle-state cadence**** (§3.5) / the ****audio thwack**** cadence ?

warm-up =
 continuous hits **without** the point-ending pause rhythm; generalizes to mid-session practice /
 between-
 game knock-ups (drop the position prior).

****You framed the consumption exactly right: "so the scorer may consume it."**** ``in_warmup`` is a
****feature**
 column******; the scorer **learns** to down-weight windows inside the span. ****Warm-up is the poster**
 child for
 no-special-casing******: the naive version is ``if t < warmup_end: drop`` ? a hand-coded branch
 hanging off a
 fuzzy boundary; a 20-second error ****deletes a real opening rally****. As a weighted feature, a
 wrong boundary
 just slightly mis-weights a couple of windows the other cues rescue. The span is **also**
 persisted as
 evidence a reconciler (§6) may use to down-weight a whole cluster, but the primary path is
 feature ? scorer.

****It is measurable today with no new labels:**** the golden labels already treat warm-up knock-ups
 / pre-serve
 as ****dead time = negatives**** (Roora starts a rally at serve-contact; fusion §1 names "warm-up
 knock-ups" as
 a precision target). So an ``in_warmup`` feature is directly A/B-able on the 6 golden videos for
****precision****
 lift ? that is the promotion gate.

3.5 Consuming a richer detector ? SHUTTLE-STATE (frame-scope producer)

A natural new ****detector****: a per-frame shuttle-state classifier ? ``in_air`` / ``racket_contact`` /
``inactive``. Note the scope: it ****processes frames ? a new per-frame signal****, like WASB, so it
 is a
****frame-scope producer in the detection layer**** ? its spec belongs in
 MULTI_SIGNAL_FUSION_PLAN.md (it is a new ``RallyCue``), ****not**** a
``WindowAnalyzer``.
 This doc records ****how the framework consumes it**** once it exists:

- a ****`WindowAnalyzer`**** turns the persisted per-frame state into window features ?
``racket_contact_count``
 (? shot cadence), ``inactive_ratio`` (? dead-time) ? that flow into the scorer as columns;
- a ****`SessionAnalyzer`**** (warm-up, §3.4) reads its cadence;
- the ****reconcilers**** (§6) get **stronger** evidence: ``trailing_dead_time`` / ``over_merged``
 currently infer
 inactivity from a velocity threshold; a learned ``inactive`` state is a more direct, occlusion-
 robust
 evidence source for "truncate / split here."

****Stage it (don't build the model first).**** Most of the state is derivable for free from signals
 we already
 have ? ``in_air`` vs ``inactive`` ? WASB velocity + visibility; ``racket_contact`` ? trajectory
****direction**
 reversals****** (fusion §4: "each reversal ? a shot") ****and**** the scaffolded ****audio thwack****
 (``backend/pipeline/audio/hit_detector.py``, ADR-007's #1 cue). So ****v0** = a derived
``shuttle_state`` analyzer******
 over existing signals (no new model); measure it. ****The dedicated per-frame vision model earns**
 its keep
 only where **v0** hits a ceiling: WASB trajectory gaps****** ? occlusion, motion blur, fast smashes ?
 where there
 is no ``(x,y)`` to threshold but an appearance model could still call ``racket_contact``. ****Label**
blocker:** we
 have ****zero**** per-frame state labels (golden labels are rally intervals); a dedicated model
 needs weak/
 distilled labels or a new (minors-gated) annotation task ? the same gate as feature-level
 fusion. If you can
 weak-label state from existing signals, prefer the derived analyzer until the gap-filling case
 is proven.

4. How signals reach the scorer ? columns, not branches

The persisted `stats` columns are exactly what the harness already lifts.

```
`load_video_candidates(...,
feature_names=[...])` (`experiment.py:429`) reads each name out of the row's `stats` JSON into a
feature
column; `predict()` (`experiment.py:203`) gates/scores on those columns and merges. So a new
producer's
signal becomes a scorer input by **adding its column name to a `Variant`'s `feature_names`** ?
no new
scoring code, no branch:
```

```
```python
```

## A learned Variant that lets the scorer weight the new signals (service-fault + warm-up

```
fusion_plus = Variant(name="fusion+analyzers",
 feature_names=[*SHUTTLE_FEATURES, "person__players_in_court",
 "service_fault__service_fault",
 "service_fault__service_fault_confidence",
 "warmup__in_warmup", "warmup__in_warmup_confidence"])
compare(golden_videos, CONTROL, fusion_plus) # LOVO-honest B ? A on the 6 golden videos
```
```

The model **learns the weights** ? down-weight a window with a high-confidence `service\_fault` or high `in\_warmup`; up-weight ?2 in-court players + multiple net-crossings. **What the scorer cannot see, it cannot special-case; what it can see, it weights from data.** Same logistic regression as ADR-004; we add columns, never a net.

5. Persistence

- **Features ? `candidate\_segments.stats`** (`database.py:240`). Namespaced `>\_\_<key>` (+ `\_confidence`). Persisted **as a shadow** even when the served decision ignores them, so any future `Variant` re-scores offline (EXPERIMENT_HARNESS §5.2). `version`/`name` recorded for provenance.
- **Spans ? window membership + a session record.** A `SessionAnalyzer`'s per-window `in\_<span>` features ride in each window's `stats` (so they re-score like any feature); the span list itself (start/end/kind/ confidence) is a small per-video record (a JSON sidecar first; a table only if it grows).
- **Events ? the `events` table** (`database.py:227`, `add\_event(segment\_id, ts, type, player, details)`). Its FK is to `play\_segments`, which exist only **after** the AI handoff confirms a candidate. So analyzer events are **buffered on the candidate** (`stats["\_events"]`) and **flushed** to `events` in Phase 4, when `\_candidate\_id` ? `segment\_id` is linked (`trajectory\_hybrid.py:283`). Candidate-level discoveries survive for offline re-scoring; the `events` table stays "events on confirmed segments."
- **Corrections record** (§6) ? report-first means corrections are **recorded, never silent.** Phase one: a JSON sidecar shaped like `experiment\_leaderboard.json` (one row per issue: segment, kind, evidence, before?after, reason, timestamp). When auto-apply is promoted, add a `corrections` table via the `PRAGMA user\_version` migration ladder (`database.py:88`, `if version < 3:`) so applied

corrections are
queryable and reversible.

6. The backward seam ? the reconciliation pass (a linter for rally decisions)

Given **predicted segments + the persisted per-frame/per-window/per-session signals**, detect where the **DECISION contradicts the EVIDENCE** and propose corrections. This is the backward dual of the analyzers:
analyzers add evidence *before* the decision; reconcilers audit the decision *against* evidence *after* it.

6.1 Contracts (proposed `backend/pipeline/reconcilers`)

```
```python
@dataclass(frozen=True)
class Issue:
 segment_index: int; kind: str # kind = rule id, e.g. "trailing_dead_time"
 evidence: Dict[str, Any] # the signal values that triggered it (auditable)
 before: Interval; after: Optional[List[Interval]] # None=drop · [iv]=corrected ·
[iv1,iv2]=split
 reason: str; confidence: float

@dataclass(frozen=True)
class Correction:
 issues: List[Issue]; corrected: List[Interval] # the FULL corrected set (idempotent)

class ConsistencyRule(ABC):
 name: str; precedence: int # FIXED ordering; lower runs first
 @abstractmethod
 def check(self, seg: Interval, signals: "SegmentSignals") -> Optional[Issue]: ...

def reconcile(segments, signals, rules) -> Correction: ... # fixed precedence; idempotent
```
```

6.2 Worked rules (fixed precedence)

```
#	Rule	Evidence contradiction	Correction
1	`trailing_dead_time`	segment still "in rally" at its end, but the shuttle went inactive earlier	truncate to the last-active frame *(the owner's misfire fix)*
2	`no_net_crossing`	`net_crossings < 1` ? nothing ever crossed the net	drop (or flag)
3	`over_merged`	a dead-time gap > merge_gap of shuttle-inactivity sits inside one segment	split at the gap
4	`boundary_slack`	first sustained shuttle activity / first net-crossing is well inside the start	tighten start to serve-contact
5	`inside_warmup`	the segment lies inside a high-confidence warmup/practice span ($3.4)	flag (drop only under apply-flag)
```

Each rule reads only **persisted evidence** (trajectory activity per frame, *net_crossings*, person occupancy, analyzer signals/events, session spans, shuttle-state) ? it is a deterministic linter, not a re-detector. As richer producers land (shuttle-state, warm-up), rules get **stronger evidence**, not new branches in the decision path.

6.3 Report-first, apply-optional (non-negotiable)

`reconcile()` returns **both** an issues report **and** a corrected set. It **NEVER** silently rewrites

(auditability + the repo's honesty/attribution working agreement). The pipeline either:

- **(default)** persists the report and **surfaces it for human review** ? production output unchanged; or
- **(flagged, `reconcile.apply`)** applies the corrected set, **recording every change** in the corrections record (before/after, rule, evidence). Rollback = flip the flag.

It is a `Variant` transformation, so its lift is A/B-measurable BEFORE auto-apply. A reconciler is a pure post-`predict` stage: `reconcile(predict(...), signals, rules).corrected`. Wrapping that as a **RECONCILED** variant lets **compare(CONTROL, RECONCILED)** report the LOVO lift on the 6 golden videos (QUALITY_ITERATIONS.md corpus) ? promote to auto-apply **only** on measured lift.

Idempotent; fixed rule precedence. **reconcile(reconcile(x)) == reconcile(x)** (truncating an already-truncated segment is a no-op); precedence is a fixed table so the result is deterministic and order-independent of input segment order.

7. No-regression discipline (identical to the harness contract)

The same four guarantees EXPERIMENT_HARNESS.md §6 makes ? applied to every producer and the reconciler:

- **Default OFF.** Analyzers gate on **analyzer_flags** (default empty); the reconciler is report-first (apply gate default off). Merged code **cannot** change served behavior until explicitly enabled.
- **Shadow-persist, then promote.** Signals/spans/issues are persisted even while ignored, so a **Variant** re-scores them offline; a signal becomes a **decision input** only when a Variant that references it wins LOVO. (EXPERIMENT_HARNESS §5.2's shadow mode.)
- **Promotion only via the eval gate.** **run_regression.py** against the golden **baseline.json** (exit 1 blocks); the experiment leaderboard records the spec hash + per-video + LOVO deltas + label-set hash.
- **Rollback = flip a config flag, never `git revert`.** The pinned **CONTROL** variant is always present.

The keystone test (extends the harness's existing invariant): **all analyzer flags OFF + reconciler report-only ? candidate rows AND predictions byte-identical to control.**

8. How it composes with fusion P2

- **Window-analyzers plug into fusion P2's enrichment hook.** **FusionSegmenter._enrich_candidate_stats** (**fusion_hybrid.py:70**) is the exact injection point §3.3 generalizes; **net_crossings** and the person features become the first registered analyzers (no behavior change). New cues are add-a-producer, per ADR-007 / MULTI_SIGNAL_FUSION_PLAN §2.
- **Session-analyzers run as a second pass** over the per-window stats the same hook persists.
- **The shuttle-state detector is a new fusion `RallyCue`** (frame-scope producer) ? it belongs in the fusion plan; this framework only **consumes** it (§3.5).
- **Reconcilers consume fusion P2's persisted signals** ? the same **candidate_segments.stats**

plus the
 persisted trajectory; no new detection.
 - Everything inherits fusion P2's control-reproducing defaults, so the fusion segmenter with all flags off
 still seeds and hands off identically to `wasb\_hybrid`.

9. Trade-offs & risks

- **Deterministic vs learned signals.** A producer may emit a verdict-shaped signal (``service_fault=1.0``), but we persist it as a **feature**, so the scorer still arbitrates. Risk: a confidently-wrong cue.
 Mitigation: the **confidence column** + the **LOVO gate** ? a bad signal earns ? 0 weight, never promotes.
- **The temporal confound is the hard core.** Warm-up-vs-play, multi-court, and held-shuttle all look identical in any single frame / frequency band ? `QUALITY_ITERATIONS`' standing lesson: "the dead-time confusion is rally-vs-non-rally activity in the same place and frequency band... only the learned model's temporal features crack it." This is *why* warm-up is a **feature the model weights**, not a heuristic span-cutter, and why session-scope structure earns its own tier.
- **Scope discipline.** Put a producer at its *natural* scope: shuttle-state is frame-scope (a detector), not a window analyzer; warm-up is session-scope, not per-window. Mis-scoping forces a producer to either recompute detection or miss the structure it needs.
- **Ordering / idempotence (reconcilers).** Fixed precedence table + an idempotence test (``reconcile?reconcile == reconcile``).
- **Signal, not verdict.** A producer that *decides* (returns keep/drop) re-introduces special-casing; reviews must reject "analyzer returns a boolean keep/drop."
- **Report-first auditability.** Reconcilers never silently mutate output; every applied change is recorded.
- **Small-N feature discipline.** Few, scale/translation-invariant columns (counts/ratios/presence, never raw coords ? `MULTI_SIGNAL_FUSION_PLAN §3.2`); the model stays a logistic regression, not a net.
- **Labels gate the learned producers.** Shuttle-state and feature-level fusion need labels we don't yet have; warm-up is measurable *now* because golden treats it as negatives. Know which producers are gated.
- **Cost & licensing.** Pure analyzers are GPU-free; person- and shuttle-state-dependent producers pull torch / a second WSL pass ? gate behind their cue flags; all deps stay MIT/Apache/BSD (ADR-002/005).
- **Event-table timing.** The buffer-then-flush wiring (§5) is the one non-obvious integration detail.

10. Generality & research positioning

This is written as a **framework**, not a badminton patch. Read sport-agnostically it is: *a registry-based, multi-scope **signal-fusion** architecture for **weakly-supervised temporal event/segment detection**, with a learned arbiter and an **auditable post-hoc correction** layer.* The pieces that generalize:

- **Multi-scope producers** ? one feature table ? a learned arbiter.
- The frame/window/session taxonomy and the "signal-not-verdict" rule are domain-independent. Any temporal problem with **weak, multi-**

modal cues
 at different time scales** (broadcast sports event spotting, surgical-phase recognition, call-center /
 meeting segmentation, activity logs) fits the same shape: add a producer, add a column, let the arbiter
 weight it, gate promotion on held-out lift.
 - **Detection ? decision (persist once, re-score forever).** Decouples the expensive, hard-to-reproduce
 perception from the cheap, deterministic decision ? the property that makes the whole thing experimentally honest and cheap to iterate.
 - **Auditable reconciliation as a first-class layer.** Most pipelines bury post-hoc fixes as silent
 if-branches; treating consistency-vs-evidence as a *report-first, A/B-measured* pass is the transferable
 methodological idea (and the human-in-the-loop / safety story).
 - **The honest measuring stick.** Interval-level, scarce ground truth + **leave-one-VIDEO-out** (group by
 match, never a held-out window from the same video) is the discipline any small-corpus temporal study
 needs to avoid flattering itself.

What is defensibly contributed vs. what is standard. Multi-modal **late fusion**, weak supervision, and
 learning a linear/logistic arbiter over hand-designed features are all standard. The *combination* offered
 here ? (a) a uniform producer-registry across **three temporal scopes**, (b) a hard **no-branch** decision
 invariant that makes extension monotone, (c) a **report-first, harness-measured reconciliation** layer, and
 (d) the **detection?decision persistence** substrate that makes every experiment a zero-risk offline
 re-score ? is what would be written up, with rally detection as the motivating hard problem.

Honesty caveat (attribution working agreement). No external results or citations are asserted here. A
 real write-up still owes: a proper **related-work survey** (temporal action segmentation/detection, sports
 event spotting, multimodal fusion, weak supervision ? to confirm which claims are actually novel), **strong
 baselines**, **per-producer ablations**, and reporting on **>1 sport / camera** to substantiate the
 generality claim. Today the corpus is **6 matches** (QUALITY_ITERATIONS) ? enough for direction-checks and
 LOVO sign-tests, **not** enough to claim a general model; corpus growth is the highest-ROI unblock. Treat
 §10 as a *positioning sketch*, not a results section.

11. Phased execution (owner-gated; ONE PR per slice, off `master`, default OFF)

- **A0 ? this doc**, indexed in `docs/README.md`, cross-linked from the peers + ADR-007. *(this deliverable; docs-only.)*
 - **A1 ? analyzer contracts + registry**
 (`WindowAnalyzer`/`AnalyzerResult`/`Signal`/`AnalyzerEvent`/
 `BoundaryHint` + `analyzer\_registry`); generalize `\_enrich\_candidate\_stats` to iterate enabled analyzers
 (none enabled ? byte-identical). Retrofit `net\_crossings` + person features as the first two analyzers.
 - **A2 ? service-fault analyzer** behind `analyzer\_flags.service\_fault` (default OFF); signal + `net\_hit`
 event + end boundary hint; shadow-persisted; offline `compare(CONTROL, +service\_fault)` LOVO.
 - **A3 ? event flush wiring** (buffer analyzer events on candidates ? `add\_event` on confirmation).

```

- **A4 ? `SessionAnalyzer` tier + `WarmupAnalyzer` v0** (leading-span heuristic over window
timings),
  default OFF; persists `warmup__in_warmup` + the span record; measured for **precision** lift
on golden.
- **A5 ? derived `shuttle_state` analyzer v0** (velocity + visibility + direction-reversals (+
audio);
  **no new model**); emits `racket_contact_count` / `inactive_ratio`; measured LOVO. *(The
dedicated
  per-frame shuttle-state DETECTOR is a separate MULTI_SIGNAL_FUSION_PLAN `RallyCue` item ?
gated on
  per-frame labels; build only if v0 is blocked by trajectory gaps.)*
- **R1 ? reconciler contracts + registry + `reconcile()`**
(`ConsistencyRule`/`Issue`/`Correction`) with
  the worked rules; pure, idempotent, unit-tested; **not** wired to apply (report-only).
- **R2 ? reconciler as a `Variant` transform**: optional post-`predict` stage; `compare(CONTROL,
RECONCILED)` LOVO; report-only in production.
- **R3 ? corrections persistence + opt-in apply**: JSON sidecar first, then the `corrections`
table
  (`user_version < 3`); apply behind `reconcile.apply`; **rollback = flip the flag.**

```

Each PR is gated by the no-regression invariant (\$7) + `run\_regression.py`. Order is a suggestion; A4/A5/R1 are independent and can be reprioritized on owner steer.

12. Verification

```

- **Offline, runnable in this dev container (no GPU/torch/cv2):** pure analyzer math on
synthetic
  trajectories + person boxes (service-fault fires on a synthetic net-halt, silent on a clean
rally;
  `WarmupAnalyzer` v0 finds a synthetic leading knock-up span; derived `shuttle_state` matches a
scripted
  velocity/visibility sequence); the **no-regression invariant** (all flags OFF + reconciler
report-only ?
  candidate rows + predictions byte-identical to control); the **reconciler idempotence** test
  (`reconcile?reconcile == reconcile`) and fixed-precedence determinism; `compare(CONTROL,
variant)`
  deterministic on synthetic candidate rows.
- **Machine-side (owner, GPU/WSL ? full pipeline can't run here):** persist analyzer signals on
a real
  clip; `compare(CONTROL, +service_fault)`, `compare(CONTROL, +warmup)`, `compare(CONTROL,
RECONCILED)`
  LOVO on the 6 golden videos; confirm the owner's **trailing-dead-time misfire** truncates, the
**held-shuttle / net-hit service feed** is down-weighted, and **warm-up windows** lose
precision-mass;
  spot-check `net_hit` events on confirmed segments.
- **Docs (A0):** links resolve; indexed in `docs/README.md`; cross-linked from
  [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md),
[EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md),
[HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md), and ADR-007.

```

97. user (2026-06-13T19:47:05.985Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\ANALYZERS_AND_RECONCILERS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

98. assistant (2026-06-13T19:47:53.984Z)

Doc rewritten as the framework-framed version. Now I'll update the README one-liner to match the new framing and stage both ? keeping the owner's working-tree README untouched via the same index-only mechanism (the README in my commit becomes `master + new line`; the working tree stays the owner's WIP).

99. user (2026-06-13T19:48:02.652Z)

```
warning: in the working copy of 'C:\Users\avidu\AppData\Local\Temp\bhi_readme_new.txt', LF will
be replaced by CRLF the next time Git touches it
warning: in the working copy of 'docs/ANALYZERS_AND_RECONCILERS.md', LF will be replaced by CRLF
the next time Git touches it
=== staged README diff vs HEAD (should be OLD analyzers line ? NEW framework line) ===
diff --git a/docs/README.md b/docs/README.md
index d3c9835..cdb3ea5 100644
--- a/docs/README.md
+++ b/docs/README.md
@@ -18,6 +18,7 @@ Live documentation for the badminton-highlight-indexer.
- [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Latest (v8) vendor labeling
guidelines mirrored from the collector's `docs/annotation/`, with the fusion-driven court-
calibration / singles-doubles additions marked `? ADDED (fusion)`.
- [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
- [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (P1 SHIPPED) for A/B + shadow + SxS
experimentation on rally-detection variants with **zero production regression** ? experiments
are decoupled from deploy (new cues default OFF, promote only on measured LOVO lift, rollback =
flip a flag). Composes the existing `backend/eval/` machinery (`calibration`, `regression`,
`metrics`); implemented in `backend/eval/experiment.py`.
+- [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated, NOT
started): a **signal-fusion framework** for rally detection ? producers emit confidence-scored
**signals** at three temporal **scopes** (frame detectors / window analyzers / session
analyzers), a learned **scorer** weights them, and an auditable **report-first reconciliation**
pass lints the decision against the evidence. The spine: **no special-casing in the decision
path** (signals?scorer, never `if/else`). Worked examples: a service-fault analyzer, a warm-
up/practice span detector, and consuming a per-frame shuttle-state detector. A peer to
`EXPERIMENT_HARNESS.md` and `MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).
- [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) ? **Living log** of rally-detection quality
iterations (hypothesis ? what shipped ? measured/expected effect ? lesson). The reverse-
chronological record; terminal-state snapshots get archived under `archives/quality-
iterations/`.
- [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated, NOT
started) for two extensibility seams: **injectable window analyzers** (forward ? pluggable per-
window cues that emit confidence-scored signals + events + boundary hints, e.g. a net-hit
service-fault detector) and a **reconciliation pass** (backward ? a report-first "linter for
rally decisions" that flags where the decision contradicts the evidence and proposes
corrections). The spine: analyzers/reconcilers emit *signals*, the scoring model *learns* their
weights ? **no special-casing in the decision path**. A peer to `EXPERIMENT_HARNESS.md` and
`MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).

=== status: staged vs unstaged (in-flight fusion files must stay unstaged) ===

M docs/ANALYZERS_AND_RECONCILERS.md
MM docs/README.md
```

100. assistant (2026-06-13T19:48:44.101Z)

The staged README now has **two** analyzers lines (NEW + OLD) ? my `git show HEAD` base already contained an analyzers line, so inserting produced a duplicate. Something's off about what HEAD points at. Let me unstage README and diagnose before touching anything further.

101. user (2026-06-13T19:48:50.307Z)

```
=== branch + recent log ===
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
cbf2236 feat(segmenters): FusionSegmenter ? shuttle + net-crossing + person-cue fusion (P2,
default-OFF) (#117)
df7d5df refactor(detectors): extract the person cue into PersonDetector (fusion P1) (#115)
ecee786 feat(eval): experiment harness ? A/B + SxS variant layer (no-regression by design)
(#114)
=== current branch ===
master
=== HEAD README: does it contain QUALITY_ITERATIONS / ANALYZERS lines? ===
```



```
HEAD README line count: 32
20: - [EXPERIMENT_HARNESS.md](EXPERIMENT_HAR
21: - [QUALITY_ITERATIONS.md](QUALITY_ITERAT
22: - [ANALYZERS_AND_RECONCILERS.md](ANALYZE
=== what did HEAD (my commit) change? ===
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
docs/README.md | 3 ++-
 1 file changed, 2 insertions(+), 1 deletion(-)
=== is master's README (df7d5df) clean of QUALITY/ANALYZERS? ===
0
```

102. assistant (2026-06-13T19:49:30.590Z)

The shared checkout was flipped to `**`master`**` by a sibling session, and the in-flight work `**merged while I was editing**`: ``#117`` (FusionSegmenter P2) and ``#118`` (README + QUALITY_ITERATIONS). My doc rewrite is now sitting as an uncommitted change `*on master*`. I need to understand what happened to my PR #116 and my doc/README on master before doing anything. Let me gather facts.

103. user (2026-06-13T19:49:39.556Z)

```
=== PR #116 status ===
Unknown JSON field: "merdgeCommit"
Available fields:
```

```
additions
assignees
author
autoMergeRequest
baseRefName
baseRefOid
body
changedFiles
closed
closedAt
closingIssuesReferences
comments
commits
createdAt
deletions
files
fullDatabaseId
headRefName
headRefOid
headRepository
headRepositoryOwner
id
isCrossRepository
isDraft
labels
latestReviews
maintainerCanModify
mergeCommit
mergeStateStatus
mergeable
mergedAt
mergedBy
milestone
number
potentialMergeCommit
projectCards
projectItems
reactionGroups
reviewDecision
reviewRequests
reviews
```

```

state
statusCheckRollup
title
updatedAt
url

=== open PRs ===
[]

=== is my doc committed on master HEAD? ===
? YES, tracked on master
=== master README lines 18-24 ===
- [ROORA_LABELING_GUIDELINES.md](ROORA_LABELING_GUIDELINES.md) ? Lates
- [EVALUATION.md](EVALUATION.md) ? How we measure and improve quality.
- [EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md) ? Design (P1 SHIPPED)
- [QUALITY_ITERATIONS.md](QUALITY_ITERATIONS.md) ? Living log of r
- [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Desig

```

Practical / ongoing

```

=== my local branch + origin refs ===

docs/analyzers-and-reconcilers 0f673ff [origin/docs/analyzers-
and-reconcilers] docs: design two rally-
detection
extensibility seams (analyzers +
reconcilers)
0f673ffe986f6258424b93e0146f3dcd4fb5b016 refs/heads/docs/analyzers-and-
reconcilers
=== full working-tree status ===
M docs/ANALYZERS_AND_RECONCILERS.md
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/arial.ttf
?? scratch/debug_wsl.py
?? scratch/keep_awake.ps1
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/stitch_min5.py
?? scratch/trim_video.py

```

104. assistant (2026-06-13T19:50:35.181Z)

Now it's clear: the shared checkout flipped to **master**, and the in-flight work merged (#117 fusion, #118 readme). Critically, **master** already has my doc + the ANALYZERS README line ? so my PR #116 appears to have merged too. My framework rewrite is sitting staged on master. Let me confirm #116's fate and whether master's doc equals my original commit, so I make a clean **new** PR rather than committing to master.

105. user (2026-06-13T19:50:42.502Z)

```

=== PR #116 definitive state ===
{"headRefName":"docs/analyzers-and-reconcilers","mergedAt":null,"number":116,"state":"CLOSED","u
rl":"https://github.com/avidullu/badminton-highlight-indexer/pull/116"}

=== which commit added my doc to master? ===
cbf2236 feat(segmenters): FusionSegmenter ? shuttle + net-crossing + person-cue fusion (P2,

```

```
default-OFF) (#117)
=== does master HEAD doc == my original commit 0f673ff doc? (no output = identical) ===
=== working-tree README modified vs HEAD? (empty = clean/identical to master) ===
=== master README ANALYZERS line (current, to be refreshed) ===
- [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated, NOT
```

106. user (2026-06-13T19:51:27.022Z)

```
22      - [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) ? Design (owner-gated,
NOT started) for two extensibility seams: **injectable window analyzers** (forward ? pluggable
per-window cues that emit confidence-scored signals + events + boundary hints, e.g. a net-hit
service-fault detector) and a **reconciliation pass** (backward ? a report-first "linter for
rally decisions" that flags where the decision contradicts the evidence and proposes
corrections). The spine: analyzers/reconcilers emit *signals*, the scoring model *learns* their
weights ? **no special-casing in the decision path**. A peer to `EXPERIMENT_HARNESS.md` and
`MULTI_SIGNAL_FUSION_PLAN.md` (ADR-007 modular rally layer).
```

107. assistant (2026-06-13T19:51:24.566Z)

Now everything's clear:

```
- **PR #116 is CLOSED, not merged.** But my original doc + README line **landed on master anyway
via #117** (the sibling's FusionSegmenter PR swept up my then-uncommitted files from the shared
tree). Master's doc is byte-identical to my original commit.
- The in-flight WIP is now all committed/merged ? **the working tree is clean** (only my staged
framework rewrite + pre-existing untracked scratch/artifacts).
```

So the path forward is now simple and clean: branch off master, commit the framework rewrite as a normal modification + refresh the README line, open a ****new**** PR (since #116 is closed). No more index surgery ? there's no in-flight WIP to protect anymore.

108. user (2026-06-13T19:51:26.337Z)

```
Switched to a new branch 'docs/analyzers-framework'
--- branch:
docs/analyzers-framework
--- staged doc present:
 1 file changed, 309 insertions(+), 136 deletions(-)
```

109. assistant (2026-06-13T19:51:37.276Z)

On the new branch with the doc staged (309 insertions / 136 deletions over master's 2-seam version). Now I'll refresh the README line to the framework framing ? a clean single-line edit (no in-flight WIP to isolate anymore).

110. user (2026-06-13T19:51:42.347Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\README.md has been updated successfully. (file state is current in your context ? no need to Read it back)

111. assistant (2026-06-13T19:51:47.945Z)

README refreshed. Let me stage it, verify only my two files are in the diff, and confirm the links now resolve (QUALITY_ITERATIONS + fusion are on master post-#117/#118).

112. user (2026-06-13T19:51:51.958Z)

```
=== staged diff stat (should be ONLY my 2 docs) ===
 docs/ANALYZERS_AND_RECONCILERS.md | 445 ++++++-----
 docs/README.md                    |    2 +-
 2 files changed, 310 insertions(+), 137 deletions(-)
=== README staged diff (should be ONLY the analyzers line changed) ===
=== cross-linked targets now exist on master? ===
  OK docs/QUALITY_ITERATIONS.md
  OK docs/MULTI_SIGNAL_FUSION_PLAN.md
```

```
OK docs/EXPERIMENT_HARNESS.md
OK docs/HOW_RALLY_DETECTION_WORKS.md
OK backend/pipeline/segmenters/fusion_hybrid.py
OK backend/pipeline/audio/hit_detector.py
```

113. assistant (2026-06-13T19:52:06.007Z)

Clean ? only my two docs are staged (the README's single bullet shows as `2 +-` = one line swapped; the grep filter mis-ate the markdown bullets, but the stat confirms it), and **all cross-link targets now exist on master** so the earlier merge-order caveat is resolved. Committing.

114. user (2026-06-13T19:52:17.407Z)

```
[docs/analyzers-framework 87abb3f] docs: extend analyzers/reconcilers into a multi-scope signal-
fusion framework
 2 files changed, 310 insertions(+), 137 deletions(-)
=== result ===
87abb3f docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework
docs/ANALYZERS_AND_RECONCILERS.md | 445 ++++++-----
docs/README.md                    |    2 +-
 2 files changed, 310 insertions(+), 137 deletions(-)
```

115. assistant (2026-06-13T19:52:20.164Z)

Committed (2 files). Now push and open the new PR ? not merging.

116. user (2026-06-13T19:52:22.695Z)

```
remote:
branch 'docs/analyzers-framework' set up to track 'origin/docs/analyzers-framework'.
To https://github.com/avidullu/badminton-highlight-indexer.git
* [new branch]      docs/analyzers-framework -> docs/analyzers-framework
```

117. user (2026-06-13T19:52:43.887Z)

<https://github.com/avidullu/badminton-highlight-indexer/pull/119>

118. assistant (2026-06-13T19:52:48.850Z)

PR opened: [#119](<https://github.com/avidullu/badminton-highlight-indexer/pull/119>). Let me verify its scope and confirm I'm not leaving the shared checkout in a surprising state.

119. user (2026-06-13T19:52:52.277Z)

```
=== PR #119 ===
{"base":"master","files":["docs/ANALYZERS_AND_RECONCILERS.md","docs/README.md"],"head":"docs/ana
lyzers-framework","state":"OPEN"}
=== current branch + clean tree (only untracked carryovers expected) ===
docs/analyzers-framework
(no tracked changes outstanding)
```

120. assistant (2026-06-13T19:53:03.986Z)

PR [#119](<https://github.com/avidullu/badminton-highlight-indexer/pull/119>) is **OPEN**, base `master`, contains exactly the two docs, and I'm on the feature branch with a clean tree (only pre-existing untracked `scratch/`/`artifacts/`). Not merged.

Let me refresh the memory checkpoint ? this turn changed the picture materially (the shared checkout flipped to master under me, #116 closed, its content landed via #117).

121. user (2026-06-13T19:53:04.678Z)

```

1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (FUSION P2 SHIPPED + collision cleanup + paper aim ? master
`35a50f3`; suite 467 green; 0 open PRs).

```

122. assistant (2026-06-13T19:53:12.510Z)

There's already a sibling-written checkpoint at the top mentioning "collision cleanup + paper aim." Let me read it so I don't duplicate or contradict it.

123. user (2026-06-13T19:53:13.119Z)

```

12     **Checkpoint:** 2026-06-14 (FUSION P2 SHIPPED + collision cleanup + paper aim ? master
`35a50f3`; suite 467 green; 0 open PRs).
13     **PR #117 MERGED = `FusionSegmenter`** (fusion P2, registry `fusion_hybrid`, **default-
OFF**): shuttle-primary (windows seed
14     from WASB/TrackNet) + adjudicating cues ? `net_crossings` (Gate A, always persisted,
GPU-free via rally_gate w/ a reused
15     axis estimate) + optional person cue (Gate B) persisting the ~5 `fusion_features.py`
features (pure-tested) via
16     `PersonDetector.detections_per_frame` over each window's ORIGINAL-video frame range.
Plugs in via **two IDENTITY hooks** on
17     `trajectory_hybrid` (`_enrich_candidate_stats`/_`_select_for_handoff`) ? wasb/tracknet
byte-identical (adversarially verified).
18     Served gating optional+default-OFF (`fusion.min_crossings`/_`min_players`) ? gates the
handoff while persisting the full
19     SUPERSET (offline substrate intact). Court mask = optional `fusion.court_polygon`.
**Design-first** (`fusion-p2-design`
20     workflow) ? **31-agent adversarial review** (25 findings?19 confirmed; headline = NO-
REGRESSION verified) ? fixes folded
21     (frame-dim guards, round() frame idx, rally_gate `estimate=` reuse, honest torch
docstrings, +6 tests). **PR #118 MERGED** =
22     README index fix (QUALITY_ITERATIONS + EXPERIMENT_HARNESS P1-SHIPPED). **? SHARED-
CHECKOUT COLLISION (lesson):** the
23     spawn_task DESIGN session ran in THIS checkout, committed `0f673ff` (the analyzers
design doc) + flipped the branch BETWEEN
24     my branch-point check and my `checkout -b` ? my fusion branch rode on top ? **#117's
squash absorbed the design doc**, so
25     **`docs/ANALYZERS_AND_RECONCILERS.md` is on master via #117** and **PR #116 was CLOSED
as redundant** (content intact on
26     master; #116's stale "OPEN do-not-merge" checkpoint below is superseded). Re-verify `git
branch` AND that no sibling commit
27     landed since, before `checkout -b`. **Owner side-aim ratified: co-author a paper**
"local, cheap & efficient rally detection
28     via multiple clever ideas" ? [[paper-coauthor-aim]]; `docs/QUALITY_ITERATIONS.md` is the
ablation-ready empirical backbone
29     (honest measured results only). **Analyzers+reconcilers DESIGN on master (owner-gated,
NOT started):** injectable per-window
30     analyzers (e.g. net-hit service-fault) + a report-first reconciliation pass; spine =
signals?learned-scorer, **no
31     special-casing** (owner directive). ?? NEXT (owner-gated): **fusion P3** (person
features ? `experiment.compare` LOVO lift on
32     the 6 golden videos ? the eager-person-compute cleanup lands here) + measure the Gate-A
held-shuttle drop on GPU; then
33     analyzers/reconcilers impl; P4 audio. Earlier this session: PR #114 harness, #115

```

person-cue extraction (see checkpoints below).

```
34
35     **Checkpoint:** 2026-06-14 (DESIGN DOC: two extensibility seams ? **PR #116 OPEN, owner-
36 gated, do NOT merge**); docs-only,
37     branch `docs/analyzers-and-reconcilers` off master `df7d5df`). Produced
38     `docs/ANALYZERS_AND_RECONCILERS.md` (367 lines, a
39     peer to EXPERIMENT_HARNESS + MULTI_SIGNAL_FUSION_PLAN) + a 1-line `docs/README.md`
40 cross-link. **Spine (owner directive):
41     NO special-casing in the decision path** ? every analyzer/reconciler emits a SIGNAL
42 (value+confidence); the SCORING MODEL
43     learns to weight it (already true in miniature: `experiment.predict`/_gate_mask` treat
44 `net_crossings`/_players_in_court`
45     as feature columns). **Seam 1 = injectable `WindowAnalyzer`s (forward)**: a registry
46 (mirrors `segmenter_registry`) of pure
47     per-window cues over the already-produced per-frame-aligned signals (WASB trajectory +
48 `PersonDetector.detections_per_frame`
49     + `rally_gate` net); each emits an `AnalyzerResult` = confidence-scored features ?
50 `candidate_segments.stats`, events ?
51     `events` table (buffer-on-candidate, flush-on-confirm), optional boundary hints; worked
52 example = a **service-fault**
53     analyzer (shuttle halts in net region ? `service_fault` signal + `net_hit` event + END
54 hint); injects at the FusionSegmenter
55     `_enrich_candidate_stats` hook, flag-gated default OFF. **Seam 2 = reconciliation pass
56 (backward)**: a report-first "linter
57     for rally decisions" ? `ConsistencyRule`s (trailing-dead-time?truncate *(owner's misfire
58 fix)*, no-net-crossing?drop,
59     over-merged?split, boundary-slack?tighten) flag decision-vs-evidence conflicts;
60 **report-first/apply-optional, NEVER silent**;
61     expressible as a harness `Variant` transform so LOVO-measurable on the 6 golden videos
62 BEFORE auto-apply; idempotent, fixed
63     precedence. ? **DESIGNED AGAINST the IN-FLIGHT fusion-P2 WIP** (uncommitted in the
64 shared tree: `fusion_hybrid.py`,
65     `fusion_features.py`, `FusionConfig.cue_flags`, `person_detector.detections_per_frame`,
66 `docs/QUALITY_ITERATIONS.md`) ? that
67     WIP was kept **byte-intact** (I isolated my single README line via git plumbing; tree is
68 CRLF / autocrlf=true, restored
```

124. user (2026-06-13T19:53:44.094Z)

The file C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

125. assistant (2026-06-13T19:53:29.725Z)

The sibling session already recorded the collision and the **paper co-authoring aim** ([[paper-coauthor-aim]] ? "local, cheap & efficient rally detection via multiple clever ideas"). That's exactly what my \$10 (generality & research positioning) now serves. Let me prepend a checkpoint for PR #119.

126. assistant (2026-06-13T19:54:11.515Z)

A system reminder shows the working-tree docs reverted to the 2-seam (master) content ? the shared checkout was flipped off my branch again by a sibling session. My framework commit is safe on the pushed branch, but let me confirm PR #119 is intact and nothing was lost.

127. user (2026-06-13T19:54:14.650Z)

```
=== current branch / HEAD (likely flipped to master by sibling) ===
feat/runlog-readability
35a50f3 docs(readme): index QUALITY_ITERATIONS + mark EXPERIMENT_HARNESS P1 SHIPPED (#118)
=== my framework commit safe on origin branch? ===
87abb3f docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework
=== PR #119 title from origin doc (must be the FRAMEWORK version) ===
```

A signal-fusion framework for rally detection ? analyzers, scopes & reconcilers

=== PR #119 status + files ===

```
{"base":"master","files":["docs/ANALYZERS_AND_RECONCILERS.md","docs/README.md"],"head":"docs/analyzers-framework","state":"OPEN"}
```

128. assistant (2026-06-13T19:54:40.326Z)

All intact and verified. Here's where things landed.

Done ? PR [#119](https://github.com/avidullu/badminton-highlight-indexer/pull/119) (open, **not merged)**

Both ideas are folded in and the doc is reframed as a general framework with a research-positioning section. Confirmed on origin: the PR branch's doc title is **"A signal-fusion framework for rally detection ? analyzers, scopes & reconcilers,"** base `master`, exactly the two doc files.

****What changed in the doc (`2-seam ? framework`, +310/?137):****

- ****§2 ? the framework:**** producers emit ****signals**** at ****three temporal scopes**** ? ***frame* detectors / *window* analyzers / *session* analyzers** ? all consumed as feature columns by a learned scorer; four invariants (signal-not-verdict · detection ? decision · report-first · promote-on-lift).
- ****§3.4 ? warm-up/practice**** as a new ****`SessionAnalyzer`**** tier (whole-timeline span ? an **`in\_warmup`** feature the scorer ****down-weights****, not a hard **`if t<warmup\_end: drop`**; measurable on golden ***today*** since warm-up is already labeled negative).
- ****§3.5 ? shuttle-state**** placed correctly as a ****frame-scope detector**** (a fusion-plan **`RallyCue`**, not a window analyzer); derive-first, dedicated-model-only-when-trajectory-gaps-bite, label blocker flagged.
- ****§10 ? generality & research positioning**** with an explicit honesty caveat (no fabricated citations/results; related-work + ablations + >1-sport still owed; n=6 = direction-check) ? the honest backbone the paper aim needs.

Important: what happened to the first PR while I was working

The shared checkout is being worked by parallel sessions, and a lot moved underneath me:

- A sibling session ****merged the fusion-P2 work (#117) and README/QUALITY updates (#118)****, and the checkout flipped to **`master`** mid-task.
- ****#117's squash absorbed my original design doc from the shared working tree**** ? so **`docs/ANALYZERS\_AND\_RECONCILERS.md`** (the 2-seam version) is ****already on master****, and ****PR #116 was closed as redundant**** (content intact, not lost). The sibling's own checkpoint records this same collision.
- So PR #119 correctly ****evolves the doc that's now on master****, and all its cross-links (QUALITY_ITERATIONS, fusion files, audio hit-detector) resolve.

Two harmless side-effects of the shared tree, for your awareness (no action needed): the local working copy has since been flipped again to a sibling's branch (**`feat/runlog-readability`**), which is why the on-disk docs momentarily show the old 2-seam text ? ****PR #119 on origin is unaffected and holds the framework version****. I left that sibling branch checked out as I found it.

Nothing merged. #119 is yours to review. I also noted the framework extension + the collision in the live memory checkpoint, and your paper aim (**`paper-coauthor-aim`**) lines up directly with the new §10.

129. user (2026-06-13T20:10:50.299Z)

can you run a research job to identify if similar patterns exist in research papers and if we are reinventing some wheel?

130. assistant (2026-06-13T20:13:18.365Z)

This is a literature-novelty assessment ? exactly the kind of structured, fan-out research task to run as a workflow, with a hard guard against the failure mode that matters most here:

****fabricated citations****. I'll decompose our framework into its distinct design patterns, run

one web-research agent per pattern, then an **adversarial verification** stage that checks every cited paper is real and hunts for closer prior art that would downgrade any novelty claim, then synthesize a "reinventing-the-wheel scorecard."

Let me author and launch it.

131. user (2026-06-13T20:14:00.722Z)

[Request interrupted by user]

132. user (2026-06-13T20:14:00.737Z)

maybe even some course corrections which we can do on this design

133. assistant (2026-06-13T20:14:25.315Z)

Good addition ? so the job should not just answer "are we reinventing the wheel?" but mine the literature for **actionable course corrections**: techniques to adopt, pitfalls to avoid, terminology to realign, and evaluation upgrades. I'll bake that into every survey agent (extract ``techniques_to_adopt` + `pitfalls``) and make a prioritized **Course Corrections** section the centerpiece of the synthesis.

Launching the workflow now (8 design-patterns × web-research ? adversarial citation/novelty verification ? synthesis). It runs in the background; I'll relay the report when it lands.

134. user (2026-06-13T20:15:51.114Z)

Workflow launched in background. Task ID: w53mgh9np

Summary: Literature novelty + course-correction assessment for the rally-detection signal-fusion framework

Transcript dir: C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\subagents\workflows\wf_1dc2acd6-052

Script file: C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\workflows\scripts\rally-framework-prior-art-wf_1dc2acd6-052.js

(Edit this file with Write/Edit and re-invoke Workflow with {scriptPath:

"C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\workflows\scripts\rally-framework-prior-art-wf_1dc2acd6-052.js"} to iterate without resending the script.)

Run ID: wf_1dc2acd6-052

To resume after editing the script: Workflow({scriptPath: "C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\workflows\scripts\rally-framework-prior-art-wf_1dc2acd6-052.js", resumeFromRunId: "wf_1dc2acd6-052"}) ? completed agents return cached results.

You will be notified when it completes. Use /workflows to watch live progress.

135. assistant (2026-06-13T20:16:08.194Z)

Understood ? I won't speculate ahead of the data. The research job (``wf_1dc2acd6-052``) is running in the background now: 8 design-patterns surveyed against the literature, each adversarially verified (real-citation check + closer-prior-art hunt), then synthesized into a memo whose centerpiece is exactly what you asked for ? a **prioritized, citation-grounded Course Corrections section** (what to change, why, expected benefit, effort, and which part of the design it touches), plus a terminology-realignment map and an honest "what's actually novel" verdict.

When it completes (I'll be notified automatically ? no need to poll), I'll:

1. Relay the memo: the reinventing-the-wheel verdict + per-pattern scorecard.
2. Walk through the concrete **course corrections grounded in the findings** ? flagging any pitfalls we're currently walking into (e.g. late-fusion calibration, over-segmentation, LOVO leakage, segmentation metrics the field expects like F1@k / edit-score / mAP@tIoU).
3. Propose which corrections to fold into the design doc, and on your go-ahead apply the agreed

ones to PR #119 (still won't merge).

You can watch live progress with `/workflows`. I'll pick this up the moment the findings land.

136. user (2026-06-13T20:19:16.043Z)

please do update the PR with your findings

137. assistant (2026-06-13T20:19:27.327Z)

The user wants the PR updated with the findings ? let me first check whether the research job has finished and pull its output.

138. user (2026-06-13T20:19:37.715Z)

<retrieval_status>not_ready</retrieval_status>

<task_id>w53mgh9np</task_id>

<task_type>local_workflow</task_type>

<status>running</status>

139. assistant (2026-06-13T20:21:34.791Z)

The research job is still running (8 patterns × survey?verify?synthesize over live web search takes a while), so the findings aren't in yet. I won't write anything into the PR until they are ? fabricating prior-art conclusions ahead of the actual results would defeat the entire point of the job (and break the honesty rule the memo itself enforces).

Here's the plan, which I'll execute automatically the moment the workflow completes (it's harness-tracked, so I'll be re-invoked on completion ? no polling needed):

1. ****Read the memo**** ? verdict + per-pattern scorecard + verified citations.
2. ****Add a new section to `docs/ANALYZERS\_AND\_RECONCILERS.md`**** on PR #119 ? a ****Prior art & course corrections**** section that:
 - states the honest reinventing-vs-novel verdict per pattern,
 - realigns our terminology to the field's (e.g. if "reconciler" ? "refinement/post-processing", "session analyzer" ? "video-level temporal context"),
 - folds the ****course corrections**** into the design (pitfalls to avoid, techniques to adopt, evaluation metrics the field expects), and
 - lists only ****verified-real**** citations (anything flagged unverifiable gets dropped, not cited).
3. ****Commit + push to the `docs/analyzers-framework` branch**** (updates PR #119) ? ****not merging****.

I'll also flag, in plain terms, anything the literature says we're currently getting wrong before I bake it in, so you can veto specific corrections.

It's still running now ? you can watch with `/workflows`. I'll pick this up and update the PR as soon as it lands.

140. user (2026-06-13T20:28:43.201Z)

<task-notification>

<task-id>w53mgh9np</task-id>

<tool-use-id>toolu_01PxKaMJ2LCLo6vVYqnP8TzC</tool-use-id>

<output-file>C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\tasks\w53mgh9np.output</output-file>

<status>completed</status>

<summary>Dynamic workflow "Literature novelty + course-correction assessment for the rally-detection signal-fusion framework" completed</summary>

<result>"I have enough to ground the terminology mapping in the actual codebase vocabulary (`RallyCue`, `WindowAnalyzer`, `ConsistencyRule`, `\_enrich\_candidate\_stats`, ADR-004 logistic

regression, `experiment.predict()`. I have everything needed to write the memo.\n\n# Rally-detection framework ? prior-art & course-correction memo\n\n## 1. Executive verdict\n\n**We are not reinventing the wheel wholesale ? but we are re-deriving, under bespoke names, a large amount of well-charted machinery, and we are currently missing several cheap, standard safeguards that the literature treats as table stakes.** Three of our five patterns are *standard and should simply be cited and reused*: the `candidate-windows + learned scorer` skeleton is textbook proposal-then-score **Temporal Action Localization** (BSN/BMN/S-CNN); the `refinement stage that fixes over-segmentation` is the defining theme of **Temporal Action Segmentation** (MS-TCN, ASRF) and **model-agnostic post-processing** (GAP, Activity Grammars); and `signal not verdict / learned arbiter over cue columns` is **late/score-level fusion** (Snoek 2005) realized as **stacked generalization** (Wolpert 1992), a pattern that already exists *in our exact domain* (racquet-sports rally/break fusion, CVIU 2009; badminton point-segmentation, Ghosh et al. 2018). The fusion mechanism itself is ~20 years old in tennis ? do not present it as new. What survives as a genuinely novel **combination** (not a research gap) is the *engineering envelope*: detection?decision as a persisted `compute-once, re-score-forever` substrate paired with a **report-first, A/B-lift-gated, idempotent, human-in-the-loop reconciler-as-linter** keyed on decision-vs-evidence *contradiction*, plus the flag-gated/default-OFF no-regression promotion harness, all on a deliberately tiny (~6-match) LOVO-by-match golden set with a deliberately interpretable logistic scorer. Of the five themes, only the **reconciler-as-linter (output-linting)** earned a `novel-combination` verdict on verification; the other four are `partial-overlap`. The two thinnest residual slices ? *warm-up/practice-phase detection as a soft session-scope feature*, and the *measured-before-apply auditable reconciliation discipline* ? are where any novelty claim should be narrowly parked.\n\n> **TL;DR:** The architecture's *skeleton* is standard TAL/TAS + late fusion (cite, don't claim); the defensible contribution is the *governance/discipline envelope* (persist-once re-score, report-first measured-before-apply reconciler, no-regression harness) applied to a scarce-label static-camera racket-sports setting ? and we are walking into known pitfalls (over-segmentation hidden by tIoU, late-fusion miscalibration, LOOCV distributional bias) that have cheap, citable fixes.\n\n## 2. Per-pattern scorecard\n\n| Pattern (our framing) | Field's name | Closest verified prior art | Verdict | One-line |\n|---|---|---|---|---|\n| Candidate-windows + learned scorer + reconciler-as-refinement | Temporal Action Localization / Segmentation + boundary/post-proc refinement | ASRF (Ishikawa et al., WACV 2021, arXiv:2007.06866); Activity Grammars (Gong et al., NeurIPS 2023, arXiv:2312.04266); BMN (Lin et al., ICCV 2019, arXiv:1907.09702) | **partial-overlap** | Our skeleton is textbook TAL/TAS; the reconciler is a hand-built activity grammar ? adopt the field's metrics & global-parse framing. |\n| `Signal not verdict`: every cue a feature column ? learned arbiter | Late/score-level fusion + stacked generalization | Snoek et al. (ACM MM 2005, 10.1145/1101149.1101236); Wolpert (Neural Networks 1992); racquet-sports A/V rally-break fusion (CVIU 2009, S1077314208001355) | **partial-overlap** | Fusion mechanism is established (and in-domain ~20 yrs); novelty is the engineering envelope, not the fusion. |\n| Detection?decision `persist once, re-score forever` | Frozen-backbone cached features + decoupled proposal/classification | S-CNN (Shou et al., CVPR 2016, arXiv:1601.02129); CBR (Gao et al., BMVC 2017, arXiv:1705.01180); DCR (Cheng/Shi et al., ECCV 2018, arXiv:1810.04002); Decouple-SSAD (arXiv:1904.07442) | **partial-overlap** | `Compute-once, re-score-many` is the *default* of the TAD field ? cite as standard practice, claim only the persist-once/idempotent discipline. |\n| Report-first reconciler (decision-vs-evidence contradiction ? auditable corrections, measured before apply) | Model assertions / consistency assertions + constraint-based data repair (`ML linter`) | Model Assertions (Kang et al., MLSys 2020, arXiv:2003.01668); HoloClean (Rekatsinas et al., VLDB 2017, arXiv:1702.00820); TFDV (Breck/Polyzotis et al., MLSys 2019); EventAnchor (CHI 2021, arXiv:2101.04954) | **novel-combination** | Ingredients all exist; the *measured-before-apply + auditable rule registry + signal-not-verdict* bundle for temporal intervals is not matched. |\n| Scarce interval labels + tiny LR + leave-one-VIDEO-out (tIoU) | Point/timestamp-supervised TAL + grouped (LOSO/LOGO) CV | SF-Net (Ma et al., ECCV 2020, arXiv:2003.06845); `Distributional bias compromises LOOCV` (Austin/Pe'er/Korem, Science Advances 2025, arXiv:2406.01652); Snorkel (Ratner et al., VLDB 2017, arXiv:1711.10160) | **partial-overlap** | Each component is named prior art; the urgent corrections are *evaluation-side* (RLOOCV, per-fold CIs), not architectural. |\n| Multi-scope (frame/window/session) hierarchy; warm-up conditions per-window | Multi-scale/hierarchical temporal modeling + play-break + phase recognition | Tjondronegoro et al. (ICME 2003 play-break); GL-MSTCN (Fang et al., 2022, 10.1186/s12938-022-01048-w); MS-TCN (Abu Farha & Gall, CVPR 2019, arXiv:1903.01945) | **partial-overlap** | Coarse-state-conditions-fine is decades old; warm-up-as-soft-feature is the only thin genuine gap. |\n| `Cue-level prior art (specific signals)` | net-crossing / in-play state machines (tennis rally detection, Sports Technology 2013, 10.1080/19346182.2013.819007); shuttle-state / hit detection (MonoTrack, Liu & Hodgins, CVPRW 2022, arXiv:2204.01899; Hsu/Yu/Cheng, Sensors 2024, 24(13):4372); service-fault (Goh et al., Sensors 2023, 23(24):9759); rally start/end on

****non-broadcast**** badminton (See et al., ISACE 2024, 10.1007/978-3-031-69073-0_15; SN Computer Science 2025, 10.1007/s42979-025-03935-0); rally/stroke dataset (ShuttleSet, KDD 2023, arXiv:2306.04948). ****All verified real.****\n\n****Unverified ? do not cite as written:**** The corpus had ****zero fabricated papers**** (every key paper resolved). But three citations carry ****wrong author lines**** that must be fixed before use (the papers are real; the attributions are not):\n- CVIU 2009 racquet-sports summarization (S1077314208001355) ? actual authors ****Liu, Huang, Jiang, Xing, Ye, Gao****, ****not**** \"Chen/Kolekar.\" Delete every \"Kolekar-line\" framing.\n- SmartTennisTV (arXiv:1801.01430) ? actual authors ****Ghosh & Jawahar****, ****not**** \"Mohan, Pramod, Jawahar.\" \n- Sensors 2024 shuttle-fusion (24(13):4372) ? actual authors ****Hsu, Yu, Cheng****, ****not**** \"Lin et al.\" \n\nAlso fix metadata slips: the L00CV-bias paper is in ****Science Advances**** (10.1126/sciadv.adx6976), not \"Briefings in Bioinformatics\"; SF-Net evaluates on ****THUMOS14/GTEA/BE0ID****, not \"ActivityNet\"; ShuttleSet (KDD) is ****arXiv:2306.04948****, while ****arXiv:2306.15664**** is the separate ShuttleSet22/stroke-forecasting paper (the two were conflated); the DeeQu entry merges the SIGMOD'19 demo with the PVLDB'18 paper.\n\n## 3. Terminology we should adopt\n\n| Our term (and where it lives in code/docs) | Field's established term | Citation anchor |\n|---|---|---|\n| \"Producer\" / \"RallyCue\" (MULTI_SIGNAL_FUSION_PLAN.md) | ****Labeling function**** / weak-supervision source; base model in a stack | Snorkel (arXiv:1711.10160); Wolpert 1992 |\n| \"Signal not verdict / no special-casing\" (ANALYZERS_AND_RECONCILERS.md §1) | ****Late / score-level fusion****; arbiter = ****level-1 combiner / meta-learner**** (stacked generalization) | Snoek et al. 2

... (truncated 25336 chars, full result in C:\\Users\\avidu\\AppData\\Local\\Temp\\claude\\C--Users-avidu-Projects-badminton-highlight-indexer\\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\\tasks\\w53mgh9np.output)</result>

<failures>[survey:sports-events] failed: API Error: The socket connection was closed unexpectedly. For more information, pass `verbose: true` in the second argument to fetch()</failures>

<usage><agent_count>16</agent_count><subagent_tokens>1074562</subagent_tokens><tool_uses>308</tool_uses><duration_ms>772090</duration_ms></usage>

</task-notification>

141. assistant (2026-06-13T20:28:57.952Z)

The research completed (16 agents; one sub-survey hit a transient socket error but the rest succeeded, and verification confirmed ****zero fabricated papers****). The notification truncated the memo ? let me read the full output, especially the Course Corrections and citation sections.

142. user (2026-06-13T20:28:59.107Z)

```

1      {
2      "summary": "Literature novelty + course-correction assessment for the rally-detection
signal-fusion framework",
3      "agentCount": 16,
4      "logs": [
5      "Surveying 8 design patterns against the literature (each: survey ? adversarial
verify)",
6      "[survey:sports-events] failed: API Error: The socket connection was closed
unexpectedly. For more information, pass `verbose: true` in the second argument to fetch()",
7      "Surveyed+verified 7/8 themes; synthesizing scorecard + course corrections"
8      ],
9      "result": "I have enough to ground the terminology mapping in the actual codebase
vocabulary (`RallyCue`, `WindowAnalyzer`, `ConsistencyRule`, `_enrich_candidate_stats`, ADR-004
logistic regression, `experiment.predict()`). I have everything needed to write the memo.\n\n#
Rally-detection framework ? prior-art & course-correction memo\n\n## 1. Executive
verdict\n\n**We are not reinventing the wheel wholesale ? but we are re-deriving, under bespoke
names, a large amount of well-charted machinery, and we are currently missing several cheap,
standard safeguards that the literature treats as table stakes.** Three of our five patterns are
*standard and should simply be cited and reused*: the \"candidate-windows + learned scorer\"
skeleton is textbook proposal-then-score **Temporal Action Localization** (BSN/BMN/S-CNN); the
\"refinement stage that fixes over-segmentation\" is the defining theme of **Temporal Action
Segmentation** (MS-TCN, ASRF) and **model-agnostic post-processing** (GAP, Activity Grammars);
and \"signal not verdict / learned arbiter over cue columns\" is **late/score-level fusion**
(Snoek 2005) realized as **stacked generalization** (Wolpert 1992), a pattern that already
exists *in our exact domain* (racquet-sports rally/break fusion, CVIU 2009; badminton point-
segmentation, Ghosh et al. 2018). The fusion mechanism itself is ~20 years old in tennis ? do
not present it as new. What survives as a genuinely novel **combination** (not a research gap)

```

is the *engineering envelope*: detection?decision as a persisted `"compute-once, re-score-forever"` substrate paired with a *report-first, A/B-lift-gated, idempotent, human-in-the-loop reconciler-as-linter* keyed on decision-vs-evidence *contradiction*, plus the flag-gated/default-OFF no-regression promotion harness, all on a deliberately tiny (~6-match) LOVO-by-match golden set with a deliberately interpretable logistic scorer. Of the five themes, only the *reconciler-as-linter (output-linting)* earned a *novel-combination* verdict on verification; the other four are *partial-overlap*. The two thinnest residual slices? *warm-up/practice-phase detection as a soft session-scope feature*, and the *measured-before-apply auditable reconciliation discipline*? are where any novelty claim should be narrowly parked.

TL;DR: The architecture's *skeleton* is standard TAL/TAS + late fusion (cite, don't claim); the defensible contribution is the *governance/discipline envelope* (persist-once re-score, report-first measured-before-apply reconciler, no-regression harness) applied to a scarce-label static-camera racket-sports setting? and we are walking into known pitfalls (over-segmentation hidden by tIoU, late-fusion miscalibration, LOOCV distributional bias) that have cheap, citable fixes.

2. Per-pattern scorecard

| Pattern (our framing) | Field's name | Closest verified prior art | Verdict |
|---|--|---|--|
| One-line | Field's name | Closest verified prior art | Verdict |
| Candidate-windows + learned scorer + reconciler-as-refinement | Temporal Action Localization / Segmentation + boundary/post-proc refinement | ASRF (Ishikawa et al., WACV 2021, arXiv:2007.06866); Activity Grammars (Gong et al., NeurIPS 2023, arXiv:2312.04266); BMN (Lin et al., ICCV 2019, arXiv:1907.09702) | <i>partial-overlap</i> |
| Our skeleton is textbook TAL/TAS; the reconciler is a hand-built activity grammar? adopt the field's metrics & global-parse framing. | Signal not verdict | every cue a feature column? learned arbiter Late/score-level fusion + stacked generalization Snoek et al. (ACM MM 2005, 10.1145/1101149.1101236); Wolpert (Neural Networks 1992); racquet-sports A/V rally-break fusion (CVIU 2009, S1077314208001355) | <i>partial-overlap</i> |
| Fusion mechanism is established (and in-domain ~20 yrs); novelty is the engineering envelope, not the fusion. | Detection?decision | "persist once, re-score forever" Frozen-backbone cached features + decoupled proposal/classification S-CNN (Shou et al., CVPR 2016, arXiv:1601.02129); CBR (Gao et al., BMVC 2017, arXiv:1705.01180); DCR (Cheng/Shi et al., ECCV 2018, arXiv:1810.04002); Decouple-SSAD (arXiv:1904.07442) | <i>partial-overlap</i> |
| "Compute-once, re-score-many" is the <i>default</i> of the TAD field? cite as standard practice, claim only the persist-once/idempotent discipline. | Report-first reconciler (decision-vs-evidence contradiction? auditable corrections, measured before apply) Model assertions / consistency assertions + constraint-based data repair ("ML linter") | Model Assertions (Kang et al., MLSys 2020, arXiv:2003.01668); HoloClean (Rekatsinas et al., VLDB 2017, arXiv:1702.00820); TFDV (Breck/Polyzotis et al., MLSys 2019); EventAnchor (CHI 2021, arXiv:2101.04954) | <i>novel-combination</i> |
| Ingredients all exist; the <i>measured-before-apply + auditable rule registry + signal-not-verdict</i> bundle for temporal intervals is not matched. | Scarce interval labels + tiny LR + leave-one-VIDEO-out (tIoU) Point/timestamp-supervised TAL + grouped (LOS0/LOG0) CV SF-Net (Ma et al., ECCV 2020, arXiv:2003.06845); "Distributional bias compromises LOOCV" (Austin/Pe'er/Korem, Science Advances 2025, arXiv:2406.01652); Snorkel (Ratner et al., VLDB 2017, arXiv:1711.10160) | <i>partial-overlap</i> | |
| Each component is named prior art; the urgent corrections are <i>evaluation-side</i> (RLOOCV, per-fold CIs), not architectural. | Multi-scope (frame/window/session) hierarchy; warm-up conditions per-window Multi-scale/hierarchical temporal modeling + play-break + phase recognition Tjondronegoro et al. (ICME 2003 play-break); GL-MSTCN (Fang et al., 2022, 10.1186/s12938-022-01048-w); MS-TCN (Abu Farha & Gall, CVPR 2019, arXiv:1903.01945) | <i>partial-overlap</i> | |
| Coarse-state-conditions-fine is decades old; warm-up-as-soft-feature is the only thin genuine gap. | Cue-level prior art (specific signals): net-crossing / in-play state machines (tennis rally detection, Sports Technology 2013, 10.1080/19346182.2013.819007); shuttle-state / hit detection (MonoTrack, Liu & Hodgins, CVPRW 2022, arXiv:2204.01899; Hsu/Yu/Cheng, Sensors 2024, 24(13):4372); service-fault (Goh et al., Sensors 2023, 23(24):9759); rally start/end on <i>non-broadcast</i> badminton (See et al., ISACE 2024, 10.1007/978-3-031-69073-0_15; SN Computer Science 2025, 10.1007/s42979-025-03935-0); rally/stroke dataset (ShuttleSet, KDD 2023, arXiv:2306.04948). | <i>All verified real.</i> | |
| Unverified? do not cite as written: The corpus had <i>zero</i> fabricated papers (every key paper resolved). But three citations carry <i>wrong author lines</i> that must be fixed before use (the papers are real; the attributions are not): | n- CVIU 2009 racquet-sports summarization (S1077314208001355)? actual authors <i>Liu, Huang, Jiang, Xing, Ye, Gao</i> , <i>not</i> "Chen/Kolekar." Delete every "Kolekar-line" framing. | SmartTennisTV (arXiv:1801.01430)? actual authors <i>Ghosh & Jawahar</i> , <i>not</i> "Mohan, Pramod, Jawahar." | n- Sensors 2024 shuttle-fusion (24(13):4372)? actual authors <i>Hsu, Yu, Cheng</i> , <i>not</i> "Lin et al." |
| Also fix metadata slips: the LOOCV-bias paper is in <i>Science Advances</i> (10.1126/sciadv.adx6976), not "Briefings in Bioinformatics"; SF-Net evaluates on <i>THUMOS14/GTEA/BE0ID</i> , not "ActivityNet"; ShuttleSet (KDD) is <i>arXiv:2306.04948</i> , while <i>arXiv:2306.15664</i> is the separate ShuttleSet22/stroke-forecasting paper (the two were conflated); the DeeQu entry merges the SIGMOD'19 demo with the PVLDB'18 paper. | | | |

3. Terminology we should adopt

Our term

(and where it lives in code/docs) | Field's established term | Citation anchor
| \n|---|---|---| \n| \n"Producer" / \nRallyCue` (MULTI_SIGNAL_FUSION_PLAN.md) | ****Labeling function**** / weak-supervision source; base model in a stack | Snorkel (arXiv:1711.10160); Wolpert 1992 | \n| \n"Signal not verdict / no special-casing" (ANALYZERS_AND_RECONCILERS.md §1) | ****Late / score-level fusion****; arbiter = ****level-1 combiner / meta-learner**** (stacked generalization) | Snoek et al. 2005; Wolpert 1992 | \n| The LR scorer (`backend/eval/classifier.py`, ADR-004) | ****Stacked meta-learner / level-1 combiner**** over producer outputs | Wolpert 1992 | \n| \n"Detection ? decision / persist once, re-score forever" | ****Frozen-backbone precomputed (cached) features****; ****decoupled proposal/classification****; offline/off-policy re-scoring | S-CNN (arXiv:1601.02129); Decouple-SSAD (arXiv:1904.07442); DOPE (arXiv:2103.16596) | \n| `WindowAnalyzer` (window-scope) | ****Temporal action proposal**** features; per-proposal ****actionness**** | S-CNN; BMN (arXiv:1907.09702) | \n| Session-scope analyzer / warm-up detection | ****Video-level temporal context / phase segmentation****; ****play-break**** outer layer; ****global-local**** modeling | Tjondronegoro ICME 2003; GL-MSTCN 2022 | \n| \n"Reconciler" (`ConsistencyRule` registry) | ****Model assertions / consistency assertions****; ****constraint-based data repair****; ****boundary refinement**** (when it tightens boundaries); ****activity grammar / constrained decoding**** (when it enforces structure) | Kang et al. MLSys 2020; HoloClean (arXiv:1702.00820); ASRF (arXiv:2007.06866); Activity Grammars (arXiv:2312.04266) | \n| \n"Report-first linter" | ****Data validation / schema anomalies**** (detect-and-report, don't silently mutate); \n"unit tests for data" | TFDV (Breck/Polyzotis et al., MLSys 2019); Deequ | \n| \n"Overmerged / split a segment" | ****Over-segmentation error**** | MS-TCN; ASRF | \n| \n"Tighten boundary slack to serve-contact" | ****Action boundary regression / boundary refinement**** | ASRF; CBR | \n| \n"Leave-one-video-out (LOVO), group by match" | ****Leave-one-group-out / leave-one-subject-out (LOGO/LOSO)**** grouped CV | scikit-learn `GroupKFold`/`LeaveOneGroupOut`; Austin et al. 2025 | \n| Temporal-IoU metric | ****mAP@tIoU**** (TAL) + ****segmental F1@{10,25,50} & edit score**** (TAS) + ****tight average-mAP / mAP@? **** (action spotting) | Lea et al. 2017; SoccerNet (arXiv:1804.04527); Soares et al. (arXiv:2206.07846) | \n\n****Concrete action:**** rename in docs/code comments so reviewers map instantly to 30 years of literature. This is a **low-effort, high-leverage framing change** and signals we know the failure modes. Keep the internal nouns (`RallyCue`, `WindowAnalyzer`, `ConsistencyRule`) as the implementation names but gloss each, on first use, with its field term. \n\n## 4. What is defensibly novel (what to claim) \n\n****CAN claim** (narrow, honest): \n1. ****The engineering/governance envelope as an integrated whole**** ? late/score-level stacked fusion **operationalized** as a no-regression, flag-gated, grouped-CV cue-promotion pipeline with a persist-once/re-score split and a TFDV-style report-first reconciler. No single found paper bundles all of: (a) report-first detection-vs-evidence **contradiction** detection over a persisted temporal decision, (b) auditable structured interval corrections (truncate/split/drop/tighten), (c) ****per-rule A/B lift-gating on golden intervals *before* auto-apply****, (d) idempotency + fixed precedence + human-in-the-loop, (e) feeding a learned (value,confidence) cue scorer. That **combination** is the contribution. \n2. ****The \n"signal not verdict" discipline as a deliberate contrast**** to the in-domain prior art (See et al. 2024/2025, Hsu et al. 2024, MonoTrack's constrained optimization, the tennis FSMs) that all ****hard-code**** these same sound cues as if/else verdicts. This is a real, defensible **design-philosophy delta** ? but state it as a philosophy delta, not a capability the prior work lacks. \n3. ****Warm-up/practice-phase detection as a soft session-scope membership feature**** ? the single slice with genuinely thin prior art (targeted searches surfaced only patents and coaching content; nothing academic does warm-up-as-soft-conditioning-feature). Claim it cautiously, graded against labels, expect to validate from scratch. \n\n****CANNOT claim** (would overstate): \n- ? The ****fusion mechanism**** (\n"multi-cue features ? light learned scorer for rally segmentation") ? ~20 years old in racquet sports (CVIU 2009; Ghosh et al. 2018 learn an SVM over per-frame HOG ? smoothed rally segments; the closest analog to our decision layer). \n- ? The ****\n"persist once, re-score forever" substrate**** as novel ? it is the **default operating mode** of the entire TAD field (THUMOS/ActivityNet ship precomputed I3D/TSN/VideoMAE features precisely so cheap heads re-train on cached features; OpenTAD/TallFormer codify Stage-0-encode/Stage-1-head). Cite as standard practice. \n- ? The ****cross-check-and-correct insight**** of the reconciler ? Hsu et al. (2024) already cross-check shuttle-vs-swing streams and correct contradictions in badminton; EventAnchor (CHI 2021) already does human-in-the-loop suggest-then-verify on racket-sports event anchors against CV evidence. Our delta is **measured-before-apply + auditability**, not the cross-check or the human-in-the-loop. \n- ? ****\n"Static single-camera / non-broadcast footage" as a novel setting**** ? Yoshikawa et al. (ceiling-cam service-scene detection) and See et al. (non-broadcast multi-player) already work there. Reframe as a robustness/engineering contribution of **our specific WASB-track + audio + person fusion on it**, not a setting novelty. \n- ? The ****no-regression / flag-gated promotion harness**** as a primitive ? textbook MLOps (shadow testing, feature-flag-gated rollout). Claim only \n"this discipline applied to a deliberately tiny golden set. \n\n\n## 5. COURSE CORRECTIONS (the actionable core) \n\nPrioritized. Each: WHAT / WHY (+cite) / BENEFIT / EFFORT / TOUCHES. \n\n\n###

Tier 1 ? Do first: these threaten the **honesty of our measured lift** on ~6 matches\n\n**C1. Switch the LOVO gate to a bias-aware protocol (RL00CV / class-balance-preserving folds; report per-fold variance + CIs, never a single mean).**\n- WHY: With ~6 match-groups, a regularized LR, and imbalanced rally/dead-time labels, standard leave-one-out induces a **negative correlation** between each fold's training-label mean and its held-out label ? models regress to the train mean, systematically biasing the estimate (a LR scored ~0.23 auROC on **random** data in the cited study). Our *"promote only on measured lift"* gate can promote noise or reject a real cue. (Austin/Pe'er/Korem, Science Advances 2025, arXiv:2406.01652.)\n- BENEFIT: Trustworthy promotion decisions; lift claims survive review.\n- EFFORT: low?med. TOUCHES: ***harness/eval***.\n\n**C2. Calibrate each producer's confidence (Platt for binary cues, isotonic for monotone ones) **before** it enters the LR.**\n- WHY: A CNN softmax, an audio onset score, and an integer net-crossing count live on wildly different scales; uncalibrated, the most over-confident cue dominates the arbiter and starves complementary cues (our audio cue's known low discrimination is a live example). (MoCaE, Oksuz et al., TMLR 2023, arXiv:2309.14976; modality competition, ICML 2022, arXiv:2203.12221.)\n- BENEFIT: Stable, meaningful weights; the *"(value,confidence)"* semantics become real rather than illusory.\n- EFFORT: med. TOUCHES: ***multi-scope producers + scorer***.\n\n**C3. Generate the arbiter's training features with OUT-OF-FOLD producer outputs (honest stacking).**\n- WHY: When a producer is itself learned (the proposed shuttle-state classifier), feeding its in-sample outputs to the LR inflates measured lift. Wolpert's stacking result; extend the group-by-match split to producer training. (Wolpert 1992.)\n- BENEFIT: Removes a leakage path that compounds with the tiny corpus.\n- EFFORT: med. TOUCHES: ***detection?decision + harness***.\n\n**C4. Report segmental F1@{10,25,50} + segmental edit score (and a tight/loose mAP@? sweep) alongside temporal-IoU.**\n- WHY: tIoU is nearly blind to over-segmentation ? a rally chopped into 3 pieces, or two rallies merged across dead-time, can still score moderate tIoU. F1@k/edit **directly** quantify the fragmentation/boundary errors our reconciler's split/merge/tighten rules exist to fix; without them a reconciler win shows no IoU lift and looks like a no-op (or an IoU-chasing rule shreds structure invisibly). Tight average-mAP (SoccerNet-style) is the metric designed for boundary precision ? directly grades the serve-contact tightening. (MS-TCN arXiv:1903.01945; ASRF arXiv:2007.06866; SoccerNet arXiv:1804.04527; Soares et al. arXiv:2206.07846.)\n- BENEFIT: Lift claims become legible/comparable to the field; rewards fixing fragmentation, not just overlap.\n- EFFORT: low. TOUCHES: ***harness/eval + reconciler***.\n\n### Tier 2 ? Architecture/discipline upgrades\n\n**C5. Frame reconciler rules as an explicit **activity grammar / constraint-aware decode** and resolve them GLOBALLY (Viterbi-style), not as stacked independent local edits.**\n- WHY: Stacked local truncate/drop/split/tighten rules with fixed precedence can conflict, oscillate, or be non-idempotent across orderings ? the exact reason the field moved to a single global parse. Activity Grammars (arXiv:2312.04266) and especially ***Constraint-Aware Decoding*** (arXiv:2605.10149, ICPR 2026 ? the **nearest** published analog to our reconciler) inject **data-derived** structural priors (transition confidence, permitted boundary sets, per-class duration) as inference-time post-processing on a frozen base and resolve via modified Viterbi. This is **our own recommended end-state already published**, and modeling priors statistically from data is a concrete template to replace hand-tuned thresholds.\n- BENEFIT: Principled, citable formalism; kills the rule-conflict/precedence headache; mitigates threshold-tuning-on-test.\n- EFFORT: med. TOUCHES: ***reconciler***.\n\n**C6. Brand the reconciler as *"model assertions / consistency assertions over the rally decision"*; attach a per-correction confidence and route only high-confidence corrections to auto-apply, low-confidence to a human queue ranked by margin.**\n- WHY: Gives a defensible lineage (Kang et al., MLSys 2020, arXiv:2003.01668 ? **and** positions us as **extending** it: they stop at flag/weak-label, we add structured correction + lift-gated apply). HoloClean (arXiv:1702.00820, ~90% repair precision ? ~10% bad repairs) and confident learning (arXiv:1911.00068) show per-repair confidence + review-most-contradicted-first is the principled way to mix auto-repair with human-in-the-loop. Treat the rule set as a TFDV-style versioned, owner-ratified schema of named anomaly types.\n- BENEFIT: Auditable, non-rotting rule registry; safer auto-apply; established vocabulary.\n- EFFORT: med. TOUCHES: ***reconciler + harness***.\n\n**C7. Add a handful of EARLY-fusion interaction features (still *"just columns,"* no branches) ? e.g. *`shuttle\_in\_air AND person\_two\_sided`* within a window.**\n- WHY: A pure-late **linear** arbiter cannot represent AND/XOR cue interactions; Snoek et al. show late usually ties early but where interactions matter early wins by a lot. Hand-crossed product columns recover this cheaply on the LR. (arXiv via 10.1145/1101149.1101236.)\n- BENEFIT: Captures cross-cue structure without leaving the *"signal not verdict"* regime.\n- EFFORT: low. TOUCHES: ***multi-scope producers + scorer***.\n\n**C8. Represent a missing/low-confidence producer output as an explicit *`is\_present`*/uncertainty companion column, not a silent 0.**\n- WHY: Collapsing *"cue absent"* and *"cue says no"* into the same 0 confuses a linear arbiter; DBF treats per-detector uncertainty as a first-class {target, non-target, unknown} state. (DBF, WACV 2016 / arXiv:1511.03183; extended arXiv:2204.02890.)\n- BENEFIT: The model learns to discount absent cues.\n- EFFORT: low. TOUCHES: ***multi-scope producers +*

scorer**. \n\n**C9. Keep the warm-up/session-scope state a SOFT membership feature column, never a hard pre-filter; add a court/context-invariance probe.** \n- WHY: A wrong coarse call in a hard gate is unrecoverable downstream (it suppresses real rallies in a mislabeled span) ? the classic hierarchical-gating risk. And **action-context confusion** (arXiv:2103.16155) is the documented trap: the static court co-occurs with rallies, so a scorer can learn \"court visible\" instead of \"point in play.\" Probe by evaluating on warm-up stretches that share the court but lack point cadence; encode the session state as a feature the LR weights (GL-MSTCN's global-local fusion + ablation discipline). \n- BENEFIT: Recoverable coarse errors; guards against the background-shortcut failure; keeps \"signal not verdict.\" \n- EFFORT: low?med. TOUCHES: **multi-scope producers (session-scope) + scorer**. \n\n### Tier 3 ? Leverage external resources & higher-effort gains \n\n**C10. Model cue CORRELATIONS in the scorer (Snorkel/MeTaL label model as a baseline; strong L2/L1, far fewer effective features than groups).** \n- WHY: Net-crossing, two-sidedness, and shuttle-in-air co-fire on the same rallies; an unregularized LR double-counts the bundle and miscalibrates. Snorkel's lesson: model source correlations or your \"independent\" evidence isn't. (arXiv:1711.10160.) EFFORT: high. TOUCHES: **scorer**. \n\n**C11. Add a temporal-smoothing pass (Kalman/HMM/Viterbi over the per-window rally posterior, or an MS-TCN-style anti-fragmentation penalty) before cutting intervals.** \n- WHY: SmartTennisTV reaches F1 ~97.5% on rally/non-rally largely by Kalman-smoothing noisy per-frame predictions into clean intervals; MS-TCN's truncated-MSE smoothing loss is the canonical over-segmentation cure. Our LR optimizes per-window labels and is blind to cross-window fragmentation. (arXiv:1801.01430; arXiv:1903.01945.) EFFORT: low?med. TOUCHES: **reconciler / scorer**. \n\n**C12. Pull ShuttleSet/ShuttleSet22 + CoachAI codebase as external eval + weak-supervision/pretraining corpus; mine a cheap weak signal (point/dead-time cadence, warm-up labels) to augment the LR.** \n- WHY: Directly attacks the ~6-match bottleneck; \"Less is More\" (arXiv:1903.00859) shows a noisy free signal + explicit noise handling can match supervised SOTA; ShuttleSet (arXiv:2306.04948) is 3,685 rallies the field recognizes. EFFORT: med?high. TOUCHES: **harness/eval + producers**. \n\n**C13. Benchmark the rule-based reconciler against a learned boundary refiner (ASRF/CBR) and ground reconciler rules in MonoTrack-style physics/cadence constraints (as soft features, not hard rules).** \n- WHY: Gives a learned baseline to *justify staying rule-based*, and MonoTrack's constrained post-processing (76.4%?89.7% hit accuracy) is validated prior art for \"use persisted trajectory + domain physics to correct raw decisions\" ? but baking \"hits ~1/sec / alternating players\" as *hard* rules would re-introduce the special-casing we forbid and break on doubles/lets/warm-up. (ASRF; CBR arXiv:1705.01180; MonoTrack arXiv:2204.01899.) EFFORT: high. TOUCHES: **reconciler**. \n\n### Pitfalls we are actively walking into (flagged for the risk register) \n- **Over-segmentation hidden by tIoU** (C4) ? THE dominant error mode of cheap segment decisions; long static-camera rallies split at occlusions, short dead-time gaps merged. A scorer optimized for per-window accuracy looks good while producing fragmented rallies. \n- **Late-fusion miscalibration / dominant-cue swamping** (C2) ? the over-confident shuttle CNN softmax or low-discrimination audio cue can dominate the LR. \n- **LOOCV distributional bias + tiny-N variance** (C1) ? promotion gate promotes noise without CIs. \n- **Stacking leakage** (C3) and **threshold-tuning-on-test** (keep *all* tunable knobs ? windowing, merge/NMS thresholds, scorer regularization, reconciler thresholds ? *inside* the LOVO loop, never on the full golden set). \n- **Proposal-recall ceiling** ? the decision stage can never recover a rally the windowing missed; measure window/proposal recall (AR@AN) separately or you'll misattribute recall losses to the scorer. \n- **Report-first decaying into silent rewrite** ? without per-rule lift gating AND a confidence gate AND idempotency tests, an over-eager rule systematically truncates correct long rallies. \n- **Static-tripod transfer gap** ? broadcast cues (overhead camera-switch, replays, crowd/commentator audio) that drive SmartTennisTV/Tjondronegoro/excitement-ranking *do not exist* on our footage; do not assume their F1s transfer or import excitement features. \n\n## 6. Risks, residual gaps & unverified claims \n- **Fabrications: none.** Across all five themes and ~40 citation checks, every key paper resolved to a real work with correct IDs. The defects are (a) the **three wrong author attributions** and several metadata slips listed in §2 ? *real papers, bad bylines/venues* ? which must be corrected before any citation, and (b) two auth-gated Springer badminton DOIs whose *full text* could not be read, though their specific quantitative claims (See et al.: shuttle-rest end-of-rally, ~88.6%/35 rallies; SN-CS 2025: BiLSTM service + ?80%-no-motion end, 187 rallies) check out from snippets. \n- **Two recent (2026) analogs verified to *exist* but not deep-read:** Constraint-Aware Decoding (arXiv:2605.10149) and TennisExpert (arXiv:2603.13397). Both are concurrent (don't void novelty) but are the nearest published precedents for, respectively, the reconciler end-state and the audio-hits?cluster?filter rally pipeline; we should read both before finalizing the related-work section. \n- **Thin-evidence claims (state plainly):** *Warm-up/practice-phase detection as a soft feature* ? searches found only patents and coaching content, no academic method; this is our strongest novelty reservation **and** our highest from-scratch-validation risk. The *measured-before-apply auditable reconciler bundle* is novel by absence-of-counterexample, not by a positive gap ? a weaker form of evidence. \n-

****Verdict where uncertain:**** the ****detection?decision substrate**** sits closest to flipping toward **"reinventing"** ? frozen-feature re-scoring is field-default, so ***only*** the persist-once/idempotent/versioned discipline is ours. Treat any **"novel substrate"** language as overclaim.\n- ****Not yet searched / open research areas:**** (1) whether a recent badminton ***full-match rally/interruption frame-level dataset*** (BWF World Tour annotations surfaced in 2025 papers) exists that would relieve scarcity and serve as an external benchmark; (2) doubles-specific cue behavior (cadence/alternation priors break); (3) calibration methods specific to ***count-valued*** cues (net-crossings) vs probabilistic ones; (4) whether RL00CV's rebalancing interacts badly with grouped (by-match) folds at N=6 ? the cited paper studies record-wise LOOCV, not group-wise.\n- ****Detector-domain-gap risk inherited downstream:**** WASB/TrackNet/MonoTrack detectors are trained on broadcast/specific rigs; per-frame shuttle accuracy degrades off-distribution (MonoTrack ~88.6%) and every cue (net-crossings, state, contact) inherits that error. ***Quantify detector accuracy on our footage before trusting any cue column*** ? currently unmeasured.\n\n### 7. Recommended citation list (verified-real only)\n\n**\$ TAL/TAS skeleton, refinement & post-processing (Theme: TAD-TAS)**\n- Temporal Action Segmentation: An Analysis of Modern Techniques ? Ding, Sener, Yao, 2022 ? arXiv:2210.10352 (TPAMI)\n- MS-TCN: Multi-Stage Temporal Convolutional Network for Action Segmentation ? Abu Farha & Gall, 2019 ? arXiv:1903.01945 (CVPR 2019)\n- ASRF: Alleviating Over-segmentation Errors by Detecting Action Boundaries ? Ishikawa et al., 2021 ? arXiv:2007.06866 (WACV 2021)\n- Activity Grammars for Temporal Action Segmentation ? Gong et al., 2023 ? arXiv:2312.04266 (NeurIPS 2023)\n- Post-Processing Temporal Action Detection (GAP) ? Nag et al., 2023 ? arXiv:2211.14924 (CVPR 2023)\n- BMN: Boundary-Matching Network ? Lin et al., 2019 ? arXiv:1907.09702 (ICCV 2019)\n- Action Spotting using Dense Detection Anchors Revisited (SoccerNet 2022) ? Soares, Shah, Biswas, 2022 ? arXiv:2206.07846\n- ***Read before citing:*** Improving TAS via Constraint-Aware Decoding ? 2026 ? arXiv:2605.10149 (closest reconciler analog)\n\n**\$ Two-stage / cascade / decoupled detection (Theme: cascade-rescore)**\n- S-CNN: Temporal Action Localization via Multi-stage CNNs ? Shou, Wang, Chang, 2016 ? arXiv:1601.02129 (CVPR 2016)\n- Cascaded Boundary Regression (CBR) ? Gao, Yang, Nevatia, 2017 ? arXiv:1705.01180 (BMVC 2017)\n- Decoupled Classification Refinement (DCR) ? Cheng, Wei, ? Shi, 2018 ? arXiv:1810.04002 (ECCV 2018)\n- Rapid Object Detection using a Boosted Cascade (Viola-Jones) ? Viola & Jones, 2001 ? 10.1109/CVPR.2001.990517\n- Decoupling Localization and Classification in SSAD ? 2019 ? arXiv:1904.07442\n\n**\$ Late/score-level fusion & stacked arbiter (Theme: late-fusion)**\n- Early versus late fusion in semantic video analysis ? Snoek, Worring, Smeulders, 2005 ? 10.1145/1101149.1101236 (ACM MM)\n- Stacked Generalization ? Wolpert, 1992 ? Neural Networks 5(2):241-259\n- MoCaE: Mixture of Calibrated Experts ? Oksuz et al., 2023 ? arXiv:2309.14976 (TMLR)\n- Modality Competition ? Huang et al., 2022 ? arXiv:2203.12221 (ICML 2022)\n- DBF: Dynamic Belief Fusion ? Lee & Kwon ? WACV 2016 / arXiv:1511.03183 (extended arXiv:2204.02890)\n- Highlight Detection with Pairwise Deep Ranking ? Yao, Mei, Rui, 2016 ? CVPR 2016 (IEEE 7780481)\n- Less is More: Learning Highlight Detection from Video Duration ? Xiong et al., 2019 ? arXiv:1903.00859 (CVPR 2019)\n- A framework for flexible summarization of racquet sports video ? ****Liu, Huang, Jiang, Xing, Ye, Gao**, 2009 ? CVIU, S1077314208001355** *(fix author line)\n\n**\$ \"ML linter\" ? assertions, repair, data validation (Theme: output-linting)**\n- Model Assertions for Monitoring and Improving ML Models ? Kang, Raghavan, Bailis, Zaharia, 2020 ? arXiv:2003.01668 (MLSys)\n- HoloClean: Holistic Data Repairs with Probabilistic Inference ? Rekatsinas et al., 2017 ? arXiv:1702.00820 (VLDB)\n- Data Validation for Machine Learning (TFDV) ? Breck, Polyzotis, Roy, Whang, Zinkevich, 2019 ? MLSys 2019\n- Snorkel: Rapid Training Data Creation with Weak Supervision ? Ratner et al., 2017 ? arXiv:1711.10160 (VLDB)\n- Confident Learning ? Northcutt, Jiang, Chuang, 2021 ? arXiv:1911.00068 (JAIR)\n- EventAnchor (racket-sports human-in-the-loop correction) ? 2021 ? arXiv:2101.04954 (CHI 2021)\n\n**\$ Scarce-label evaluation, point-supervision, CV pitfalls (Theme: weak-sup-eval)**\n- SF-Net: Single-Frame Supervision for TAL ? Ma et al., 2020 ? arXiv:2003.06845 (ECCV 2020; eval on THUMOS14/GTEA/BE0ID)\n- Point-Level Temporal Action Localization ? Ju, Zhao et al., 2020 ? arXiv:2012.08236\n- Distributional bias compromises leave-one-out cross-validation ? Austin, Pe'er, Korem, 2025 ? arXiv:2406.01652 / 10.1126/sciadv.adx6976 (Science Advances)\n- WTAL via Explicit Subspaces (action-context confusion) ? Liu et al., 2021 ? arXiv:2103.16155 (AAAI 2021)\n- SoccerNet (dataset + tolerance mAP@?) ? Giancola et al., 2018 ? arXiv:1804.04527\n\n**\$ Hierarchical / play-break / phase (Theme: multiscale-phase)**\n- Generic play-break event detection for hierarchical sports video ? Tjondronegoro, Chen, Pham, 2003 ? IEEE ICME 2003\n- GL-MSTCN (global-local phase recognition) ? Fang et al., 2022 ? 10.1186/s12938-022-01048-w\n- PCL-Former (hierarchical context-aware TAL) ? 2025 ? arXiv:2507.06411 *(authorship unverified ? cite without asserting authors)\n\n**\$ Domain ? badminton/racket-sport rally & cue prior art (Themes: cue-novelty + cross-cutting)**\n- Towards Structured Analysis of Broadcast Badminton Videos (point segmentation) ? Ghosh, Singh, Jawahar, 2018 ? arXiv:1712.08714 (WACV 2018)\n- SmartTennisTV ? ****Ghosh & Jawahar**, 2018 ? arXiv:1801.01430** *(fix author line)\n- MonoTrack ? Liu & Hodgins et al., 2022 ? arXiv:2204.01899 (CVPRW 2022)\n- Enhancing Badminton Game Analysis:

Shot Refinement via Fusion ? **Hsu, Yu, Cheng**, 2024 ? Sensors 24(13):4372, 10.3390/s24134372
 (fix author line)\n- Service and End of Rally Detection in Badminton Videos ? See et al., 2024
 ? 10.1007/978-3-031-69073-0_15 (closest non-broadcast head-to-head baseline)\n- A Comparative
 Analysis ? Rally Detection in Badminton ? See et al., 2025 ? 10.1007/s42979-025-03935-0\n-
 Automated Service Height Fault Detection ? Goh et al., 2023 ? Sensors 23(24):9759\n- Automatic
 rally detection on broadcast tennis videos ? 2013 ? 10.1080/19346182.2013.819007\n- ShuttleSet ?
 Wang, Huang et al., 2023 ? arXiv:2306.04948 / 10.1145/3580305.3599906 (KDD 2023); ShuttleSet22 ?
 arXiv:2306.15664\n- Unsupervised Temporal Feature Aggregation (unstructured table-tennis rally
 detection) ? Chaudhury et al., 2019 ? arXiv:2002.08097 (ISM 2019)\n\n--\n\n**Key files
 referenced (absolute paths):** `C:\\Users\\avidu\\Projects\\badminton-highlight-
 indexer\\docs\\ANALYZERS_AND_RECONCILERS.md` (the reconciler/analyzer seam spec ? §1 \"signal
 not verdict\" and the DETECTION/DECISION/RECONCILER diagram this memo maps to the field's
 vocabulary), `docs\\MULTI\_SIGNAL\_FUSION\_PLAN.md` (`RallyCue` producers,
 `enrich\_candidate\_stats`, scale/translation-invariant features), `docs\\EXPERIMENT\_HARNESS.md`
 (`Variant`/`compare`/`compare\_unlabeled`, the no-regression A/B + LOVO gate that C1/C3/C4
 touch), `docs\\HOW\_RALLY\_DETECTION\_WORKS.md` (2-layer mental model),
 `backend\\eval\\classifier.py` (the ADR-004 logistic-regression scorer = the \"level-1
 combiner\" that C2/C7/C8/C10 modify), and `backend\\eval\\experiment.py:203`
 (`experiment.predict()` ? `net\_crossings`/`players\_in\_court` already treated as feature
 columns)."
 10 }

143. assistant (2026-06-13T20:29:25.114Z)

Full memo retrieved ? it's thorough, citation-verified (zero fabrications; a few author-line
 fixes flagged), and gives exactly the prioritized course corrections you asked for. Before
 editing, I must get back onto my PR branch (the shared checkout was flipped to a sibling's
 branch earlier) without disturbing anyone's WIP.

144. user (2026-06-13T20:29:27.811Z)

```
=== current branch + status ===
feat/fusion-p4-audio
  (no tracked changes / clean but for untracked carryovers)
=== is docs/analyzers-framework still at my framework commit on origin? ===
87abb3f docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework
```

145. assistant (2026-06-13T20:30:16.964Z)

The checkout is now on yet another sibling branch (`feat/fusion-p4-audio`, clean tree). Rather
 than flip the shared checkout again and risk stepping on the active sibling, I'll use a
 dedicated worktree for my branch ? that leaves the main checkout untouched.

146. user (2026-06-13T20:30:19.621Z)

```
Preparing worktree (checking out 'docs/analyzers-framework')
HEAD is now at 87abb3f docs: extend analyzers/reconcilers into a multi-scope signal-fusion
framework
--- worktree HEAD:
87abb3f docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework
--- doc title in worktree:
```

A signal-fusion framework for rally detection ? analyzers, scopes & reconcilers

```
--- doc line count:
540 C:/Users/avidu/Projects/bhi-analyzers-wt/docs/ANALYZERS_AND_RECONCILERS.md
```

147. assistant (2026-06-13T20:30:24.918Z)

Worktree ready with my framework doc. Let me read its tail (§10?§12) to find exact anchor points
 for the additions.

148. user (2026-06-13T20:30:25.571Z)

474 methodological idea (and the human-in-the-loop / safety story).

```

475     - **The honest measuring stick.** Interval-level, scarce ground truth + **leave-one-
VIDEO-out** (group by
476     match, never a held-out window from the same video) is the discipline any small-corpus
temporal study
477     needs to avoid flattering itself.
478
479     **What is defensibly contributed vs. what is standard.** Multi-modal **late fusion**,
weak supervision, and
480     learning a linear/logistic arbiter over hand-designed features are all standard. The
*combination* offered
481     here ? (a) a uniform producer-registry across **three temporal scopes**, (b) a hard
**no-branch** decision
482     invariant that makes extension monotone, (c) a **report-first, harness-measured
reconciliation** layer, and
483     (d) the **detection?decision persistence** substrate that makes every experiment a zero-
risk offline
484     re-score ? is what would be written up, with rally detection as the motivating hard
problem.
485
486     **Honesty caveat (attribution working agreement).** No external results or citations are
asserted here. A
487     real write-up still owes: a proper **related-work survey** (temporal action
segmentation/detection, sports
488     event spotting, multimodal fusion, weak supervision ? to confirm which claims are
actually novel), **strong
489     baselines**, **per-producer ablations**, and reporting on **>1 sport / camera** to
substantiate the
490     generality claim. Today the corpus is **6 matches** (QUALITY_ITERATIONS) ? enough for
direction-checks and
491     LOVO sign-tests, **not** enough to claim a general model; corpus growth is the highest-
ROI unblock. Treat
492     §10 as a *positioning sketch*, not a results section.
493
494     ---
495
496     ## 11. Phased execution (owner-gated; ONE PR per slice, off `master`, default OFF)
497
498     - **A0 ? this doc**, indexed in `docs/README.md`, cross-linked from the peers + ADR-007.
*(this
499     deliverable; docs-only.)*
500     - **A1 ? analyzer contracts + registry**
(`WindowAnalyzer`/`AnalyzerResult`/`Signal`/`AnalyzerEvent`/
501     `BoundaryHint` + `analyzer_registry`); generalize `_enrich_candidate_stats` to iterate
enabled analyzers
502     (none enabled ? byte-identical). Retrofit `net_crossings` + person features as the
first two analyzers.
503     - **A2 ? service-fault analyzer** behind `analyzer_flags.service_fault` (default OFF);
signal + `net_hit`
504     event + end boundary hint; shadow-persisted; offline `compare(CONTROL,
+service_fault)` LOVO.
505     - **A3 ? event flush wiring** (buffer analyzer events on candidates ? `add_event` on
confirmation).
506     - **A4 ? `SessionAnalyzer` tier + `WarmupAnalyzer` v0** (leading-span heuristic over
window timings),
507     default OFF; persists `warmup_in_warmup` + the span record; measured for
**precision** lift on golden.
508     - **A5 ? derived `shuttle_state` analyzer v0** (velocity + visibility + direction-
reversals (+ audio);
509     **no new model**); emits `racket_contact_count` / `inactive_ratio`; measured LOVO.
*(The dedicated
510     per-frame shuttle-state DETECTOR is a separate MULTI_SIGNAL_FUSION_PLAN `RallyCue`
item ? gated on
511     per-frame labels; build only if v0 is blocked by trajectory gaps.)*
512     - **R1 ? reconciler contracts + registry + `reconcile()`**
(`ConsistencyRule`/`Issue`/`Correction`) with

```

```

513     the worked rules; pure, idempotent, unit-tested; **not** wired to apply (report-only).
514 - **R2 ? reconciler as a `Variant` transform**: optional post-`predict` stage;
`compare(CONTROL,
515     RECONCILED)` LOVO; report-only in production.
516 - **R3 ? corrections persistence + opt-in apply**: JSON sidecar first, then the
`corrections` table
517     (`user_version < 3`); apply behind `reconcile.apply`; **rollback = flip the flag.**
518
519     Each PR is gated by the no-regression invariant (§7) + `run_regression.py`. Order is a
suggestion; A4/A5/R1
520     are independent and can be reprioritized on owner steer.
521
522     ---
523
524     ## 12. Verification
525
526 - **Offline, runnable in this dev container (no GPU/torch/cv2):** pure analyzer math on
synthetic
527     trajectories + person boxes (service-fault fires on a synthetic net-halt, silent on a
clean rally;
528     `WarmupAnalyzer` v0 finds a synthetic leading knock-up span; derived `shuttle_state`
matches a scripted
529     velocity/visibility sequence); the **no-regression invariant** (all flags OFF +
reconciler report-only ?
530     candidate rows + predictions byte-identical to control); the **reconciler
idempotence** test
531     (`reconcile?reconcile == reconcile`) and fixed-precedence determinism;
`compare(CONTROL, variant)`
532     deterministic on synthetic candidate rows.
533 - **Machine-side (owner, GPU/WSL ? full pipeline can't run here):** persist analyzer
signals on a real
534     clip; `compare(CONTROL, +service_fault)`, `compare(CONTROL, +warmup)`,
`compare(CONTROL, RECONCILED)`
535     LOVO on the 6 golden videos; confirm the owner's **trailing-dead-time misfire**
truncates, the
536     **held-shuttle / net-hit service feed** is down-weighted, and **warm-up windows** lose
precision-mass;
537     spot-check `net_hit` events on confirmed segments.
538 - **Docs (A0):** links resolve; indexed in `docs/README.md`; cross-linked from
539     [MULTI_SIGNAL_FUSION_PLAN.md](MULTI_SIGNAL_FUSION_PLAN.md),
[EXPERIMENT_HARNESS.md](EXPERIMENT_HARNESS.md),
540     [HOW_RALLY_DETECTION_WORKS.md](HOW_RALLY_DETECTION_WORKS.md), and ADR-007.

```

149. assistant (2026-06-13T20:31:04.431Z)

Now I'll fold the findings in. First, append a comprehensive new section (§13) with the verdict, terminology realignment, course corrections, pitfalls, and verified citations (applying the author-line/metadata fixes the verification flagged). Then two small consistency edits to §10 and §12.

150. user (2026-06-13T20:32:17.671Z)

The file C:/Users/avidu/Projects/bhi-analyzers-wt/docs/ANALYZERS_AND_RECONCILERS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

151. assistant (2026-06-13T20:32:23.224Z)

Now two consistency edits: point §10 forward to the completed review, and add the field-standard metrics to §12.

152. user (2026-06-13T20:32:28.813Z)

The file C:/Users/avidu/Projects/bhi-analyzers-wt/docs/ANALYZERS_AND_RECONCILERS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

153. user (2026-06-13T20:32:39.983Z)

The file C:/Users/avidu/Projects/bhi-analyzers-wt/docs/ANALYZERS_AND_RECONCILERS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

154. assistant (2026-06-13T20:32:47.916Z)

Now commit from the worktree and push to update PR #119 (not merging), then clean up the worktree.

155. user (2026-06-13T20:33:03.222Z)

```
=== diff stat (should be ONLY the one doc) ===
 docs/ANALYZERS_AND_RECONCILERS.md | 176 ++++++
 1 file changed, 175 insertions(+), 1 deletion(-)
=== line count now ===
594
[docs/analyzers-framework 0b6924b] docs: ground the framework in a verified literature review +
course corrections
 1 file changed, 175 insertions(+), 1 deletion(-)
=== push ===
To https://github.com/avidullu/badminton-highlight-indexer.git
 87abb3f..0b6924b docs/analyzers-framework -> docs/analyzers-framework
```

156. assistant (2026-06-13T20:33:08.312Z)

Pushed ? PR #119 updated. Now clean up the worktree and verify the PR + that the main checkout (sibling's branch) is undisturbed.

157. user (2026-06-13T20:33:11.898Z)

```
worktree removed
=== worktrees now ===
C:/Users/avidu/Projects/badminton-highlight-indexer 53d4d1b [feat/fusion-p4-audio]
=== main checkout still on sibling branch, untouched? ===
feat/fusion-p4-audio
=== PR #119 commits ===
{"commits":["docs: extend analyzers/reconcilers into a multi-scope signal-fusion f?","docs:
ground the framework in a verified literature review + course c?"],"state":"OPEN"}
```

158. assistant (2026-06-13T20:33:35.699Z)

Done ? PR [#119](https://github.com/avidullu/badminton-highlight-indexer/pull/119) is updated with the literature review (new commit `0b6924b`, **not merged**). The worktree is removed and the main checkout is untouched on the sibling's `feat/fusion-p4-audio`.

What the research found

****Are we reinventing the wheel? Partly ? and that's fine, as long as we cite honestly.**** The job surveyed 8 design patterns, adversarially verified ~40 citations (****zero fabrications****), and landed a clear verdict:

- ****The skeleton is standard ? cite, don't claim:****
 - candidate-windows + learned scorer = ****Temporal Action Localization**** (proposal-then-score: BMN, S-CNN)
 - reconciler-fixing-fragmentation = ****Temporal Action Segmentation**** refinement (MS-TCN, ASRF)
 - "signal not verdict / learned arbiter over cue columns" = ****late/score-level fusion + stacked generalization**** (Snoek 2005, Wolpert 1992) ? and this ****already exists in racket sports* (~20 years old; Ghosh et al. 2018 do almost exactly our decision layer).**
- ****Only one pattern rated `novel-combination`:**** the ****report-first reconciler-as-linter**** bundle (contradiction-detection + auditable structured corrections + ****per-rule A/B lift-gating before auto-apply* + idempotency + human-in-the-loop).** The other thin novelty is ****warm-up-as-a-soft-session-feature**** (almost no academic prior art).
- ****Defensible claim = the governance/discipline envelope****, not the architecture or the cues.

The course corrections I folded in (§13, owner-gated)

Tier 1 ? these protect the honesty of our "measured lift" on 6 matches (the ones I'd act on first):

- **C1** ? our LOVO gate is biased at N=6: plain leave-one-out anti-correlates train/held-out label means (a LR scored ~0.23 AUC on *random* data in the cited study). Switch to balance-preserving folds + **report per-fold CIs, never a single mean**.
- **C2** ? calibrate each producer's confidence before the LR, or the over-confident shuttle CNN / weak audio cue swamps the arbiter (our audio cue's low discrimination is a live example).
- **C3** ? out-of-fold stacking for the learned shuttle-state producer (else leakage inflates lift).
- **C4** ? report **F1@{10,25,50}** + edit score + **mAP@tIoU**; plain temporal-IoU is *blind to over-segmentation*, so a reconciler win (or a rule shredding rallies) is invisible under our current metric.

Tier 2/3 add: resolve reconciler rules **globally** (activity-grammar/Viterbi, not stacked local edits ? fixes the precedence/idempotence headache); brand the reconciler as "model assertions"; early-fusion interaction columns; explicit ``is_present`` columns; keep warm-up a **soft** feature + a court-context-invariance probe (the static court is a shortcut trap); ShuttleSet as external eval to attack the 6-match bottleneck.

Honesty fixes applied to the citations: three real papers had wrong author-lines in the first draft (corrected: CVIU 2009 = Liu/Huang/?, SmartTennisTV = Ghosh & Jawahar, Sensors 2024 = Hsu/Yu/Cheng), plus a few metadata slips. Two 2026 analogs (arXiv:2605.10149, arXiv:2603.13397) are flagged **"read before citing."**

If you want, I can now turn the Tier-1 corrections (C1-C4) into concrete harness changes ? but those touch ``backend/eval/`` code, so that'd be a separate implementation PR off master, owner-gated, not this docs one.

159. user (2026-06-13T20:33:38.108Z)

can I also get a summary of findings and its impacts

160. assistant (2026-06-13T20:34:25.592Z)

Summary of findings & their impacts

The one-line verdict

We didn't reinvent the wheel *wholesale*, but the **architecture skeleton** is standard prior art (cite it, don't claim it), and our **measurement methodology** has honesty bugs that could make the whole "promote-on-measured-lift" discipline lie to us. The genuinely defensible contribution is narrow: the **governance envelope** + the **report-first reconciler-as-linter** + **warm-up-as-a-soft-feature**.

Findings

| What we built | What the literature says | Verdict |
|---|--|--------------------------|
| candidate-windows + learned scorer | Temporal Action Localization (BMN, S-CNN) | standard |
| reconciler fixing merges/splits | Temporal Action Segmentation refinement (MS-TCN, ASRF) | standard |
| "signal not verdict" / cue-columns ? arbiter | late/score-level fusion + stacked generalization (Snoek '05, Wolpert '92) ? *already in racket sports | standard (~20 yrs old) |
| detection ? decision, persist-and-re-score | frozen-feature re-scoring = the default of the TAD field | standard |
| scarce labels + tiny LR + LOVO | point-supervised TAL + grouped CV | standard components |
| report-first reconciler bundle (contradiction-detect + audited correction + per-rule lift-gating + idempotent + human-in-loop) | ingredients exist (model assertions, HoloClean, TFDV) but not bundled | novel combination |
| warm-up as a soft session feature | only patents/coaching content found | thin genuine gap |

Verification was clean: **zero fabricated papers** across ~40 checks (a few author-

lines/metadata were wrong and I corrected them).

Impacts ? what actually changes because of this

****1. Evaluation honesty (highest stakes ? our measurement may be misleading us).****

This is the impact that matters most, because everything downstream rests on "did the cue measurably help?"

- ****LOVO at N?6 is biased**** ? plain leave-one-out anti-correlates train/held-out label means; a logistic regression scored ~0.23 AUC on **random** data in the cited study. ? Our gate can ****promote noise or reject a real cue****. Fix: balance-preserving folds + per-fold CIs (never a single mean).
- ****Temporal-IoU is blind to over-segmentation**** ? a rally chopped in 3, or two rallies merged, still scores OK on tIoU. ? A reconciler win is ****invisible****, and an IoU-chasing rule that shreds rallies looks fine. Fix: report F1@k + edit-score + mAP@tIoU.
- Plus leakage (in-sample learned-producer outputs) and threshold-tuning-on-test. ? ****Net impact: our "no-regression / promote-on-lift" guarantee is only as honest as these fixes. They're cheap and should land before we trust any cue measurement.****

****2. What we can claim (paper / positioning).****

- We ****cannot**** claim the fusion mechanism, the persist-and-re-score substrate, the static-camera setting, or the MLOps harness as novel ? a reviewer would reject those instantly.
- We ****can**** claim the integrated governance envelope, the "no-branch" design philosophy as a deliberate delta vs in-domain work that hard-codes cues, and warm-up-as-soft-feature.
- ? ****Impact: the paper's "contributions" and "related work" sections get rewritten now****, before we invest in experiments, so we don't build a story a reviewer kills.

****3. Design robustness (silent-failure prevention).****

Calibrate producers (or the over-confident shuttle CNN swamps the weak audio cue); ``is_present`` columns (so "cue absent" ? "cue says no"); keep warm-up a **soft** feature + a court-context-invariance probe (the static court is a shortcut the scorer can learn instead of "point in play"); resolve reconciler rules **globally** (Viterbi) instead of stacked local edits that conflict. ? ****Impact: these prevent real, hard-to-spot correctness failures****, not just cosmetics.

****4. Positive validation (de-risks the direction).****

The literature **confirms** several instincts were right: late fusion is the correct frame; report-first auditable correction has real lineage (model assertions); soft-feature-not-hard-gate is exactly what the field would advise; the no-branch discipline is a meaningful philosophy delta. ? ****Impact: we're building on established ground with the right vocabulary, not off in the weeds ? and §13 gives us that vocabulary so reviewers map us to 30 years of work.****

****5. Newly-surfaced risks (now in the doc's risk register).****

Detector domain-gap is ****unmeasured**** (WASB/TrackNet accuracy on **our** footage is unknown, and every cue inherits that error); the windowing stage sets a recall ceiling the scorer can't recover; broadcast-derived baselines (SmartTennisTV F1 ~97.5%) ****don't transfer**** to our static-tripod footage.

Net takeaway

The framework's **bones** are sound and well-precedented ? so the win isn't novelty, it's ****rigor****: adopt the field's metrics + evaluation safeguards (Tier-1 corrections) so our measured-lift claims are trustworthy, narrow the novelty claims to what survives scrutiny, and rename to the field's terms. The single most important consequence: ****fix C1+C4 before believing any cue-helps-or-not result the harness currently produces.****

If useful, I can turn C1?C4 into a concrete ``backend/eval/`` implementation plan (separate owner-gated PR off master).

161. user (2026-06-13T20:36:06.615Z)

please fork that work in a separate branch and merge the PRs with updated docs and CI

162. assistant (2026-06-13T20:37:59.339Z)

Before I merge anything, let me confirm the current PR landscape and CI status ? and I want to

flag one thing up front: merging the **docs** PR is low-risk and you've clearly asked for it, but the **course-correction code** (C1?C4) is a substantial change to our **measurement methodology** that arguably wants review before it lands on master. Let me gather facts first.

163. user (2026-06-13T20:38:04.885Z)

=== open PRs ===

119 docs/analyzers-framework [@avidullu] docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework (warm-up + shuttle-state)

=== #119 mergeable + CI checks ===

```
{"mergeStateStatus":"CLEAN","mergeable":"MERGEABLE","reviewDecision":""}
Full test suite (offline, mocked)      pass    1m32s  https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478314263/job/81221518068
Mypy (lenient baseline) pass          28s      https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478314263/job/81221518036
Ruff lint      pass          9s      https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478314263/job/81221518045
Syntax sweep + import smoke (no GPU stack) pass    24s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478314263/job/81221518047
```

164. assistant (2026-06-13T20:38:47.806Z)

PR #119 is green on all four CI lanes (full suite, mypy, ruff, import-smoke), MERGEABLE, and it's the only open PR ? and it's docs-only, which is squarely within my merge authority and what you asked for. Merging it now.

165. user (2026-06-13T20:38:56.622Z)

=== confirm merged + master HEAD ===

```
{"mergedAt":"2026-06-13T20:38:51Z","state":"MERGED"}
23672f2 docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework (warm-up + shuttle-state) (#119)
```

166. assistant (2026-06-13T20:39:55.674Z)

Merged ? PR #119 is now on master (`23672f2`); the framework doc + the literature-review §13 (verdict, terminology, course corrections, verified citations) are landed.

Now the course-correction work. One genuine fork I want your call on before I proceed, because it's different in kind from the docs: **C1?C4 change our measurement methodology** (how we compute and trust "measured lift"), so a buggy implementation could quietly distort every future cue evaluation. I'll build it as **additive / no-regression** (new functions + opt-in, never silently changing existing `compare`/LOVO defaults or baselines), on a fresh branch off master, with offline unit tests + doc + CI ? but how it lands is your decision:

167. user (2026-06-13T20:40:21.546Z)

Your questions have been answered: "How should I land the C1?C4 measurement-honesty course corrections (bias-aware grouped CV + per-fold CIs, producer calibration, out-of-fold stacking, F1@k/edit-score/mAP@tIoU metrics)?"="Implement + self-merge on green CI". You can now continue with these answers in mind.

168. assistant (2026-06-13T20:41:40.384Z)

You chose self-merge on green CI ? I'll implement C1?C4 as **additive, no-regression** helpers (new modules + tests; existing `compare`/LOVO behavior unchanged), open a PR, and merge once all CI lanes pass. To avoid the shared-checkout collisions, I'll work in a fresh worktree off `master`.

First, let me set up the worktree and read the exact `backend/eval/` APIs I'm extending (plus the CI/test runner and whether numpy is available locally).

169. user (2026-06-13T20:41:45.649Z)

```

branch 'feat/eval-honesty-corrections' set up to track 'origin/master'.
HEAD is now at 23672f2 docs: extend analyzers/reconcilers into a multi-scope signal-fusion
framework (warm-up + shuttle-state) (#119)
=== HEAD ===
23672f2 docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework (warm-up +
shuttle-state) (#119)
=== eval modules present ===
__init__.py annotations.py audio batch.py calibrate_local.py calibrate_wasb.py calibration.py
classifier.py distill_local.py eval_partial_labels.py experiment.py export_wasb_segments.py
fusion_audio.py fusion_compare.py fusion_golden.py gemini_refine.py gt_loader.py harness.py
metrics.py rally_seq_proto.py regression.py training.py
=== numpy available locally? ===
numpy 2.4.6
=== test runner / CI ===
C:/Users/avidu/Projects/bhi-eval-wt/run_all_tests.py
ci.yml

```

170. assistant (2026-06-13T20:41:54.395Z)

numpy is available locally, so I can run tests here before pushing. Let me read the modules I'm extending, the test runner, and CI to match conventions exactly.

171. user (2026-06-13T20:41:54.903Z)

```

1      """Pure temporal interval-matching metrics for rally evaluation.
2
3      No I/O, no DB, no global state ? every function here is deterministic and unit
4      tested. An "interval" is a `(start, end)` tuple of floats in seconds with
5      start < end. We score a list of *predicted* rally intervals against a list of
6      *ground-truth* intervals via temporal IoU and one-to-one greedy matching.
7      """
8
9      from dataclasses import dataclass, field
10     from typing import List, Tuple, Optional
11
12     Interval = Tuple[float, float]
13
14
15     def interval_iou(a: Interval, b: Interval) -> float:
16         """Temporal intersection-over-union of two intervals.
17
18         Returns 0.0 when they don't overlap (or either is degenerate). IoU is
19         overlap / union, both measured as durations on the timeline.
20         """
21         a_start, a_end = a
22         b_start, b_end = b
23         if a_end <= a_start or b_end <= b_start:
24             return 0.0
25         inter = max(0.0, min(a_end, b_end) - max(a_start, b_start))
26         if inter <= 0.0:
27             return 0.0
28         union = (a_end - a_start) + (b_end - b_start) - inter
29         return inter / union if union > 0 else 0.0
30
31
32     @dataclass
33     class Match:
34         """One matched predicted?ground-truth pair."""
35         pred_index: int
36         gt_index: int
37         iou: float
38
39
40     @dataclass
41     class MatchResult:

```



```

42     """Outcome of matching predictions to ground truth at a given IoU threshold."""
43     matches: List[Match] = field(default_factory=list)
44     unmatched_pred_indices: List[int] = field(default_factory=list) # false positives
45     unmatched_gt_indices: List[int] = field(default_factory=list)   # false negatives
46 (misses)
47
48     def match_intervals(preds: List[Interval], gts: List[Interval],
49                         iou_threshold: float = 0.5) -> MatchResult:
50         """Greedy one-to-one matching of predictions to ground truth by descending IoU.
51
52         Every candidate (pred, gt) pair with IoU >= threshold is considered; pairs
53         are claimed highest-IoU first, and each prediction/GT can be used once. This
54         keeps a single prediction from being credited for two ground-truth rallies
55         (and vice versa).
56         """
57         candidates = []
58         for pi, p in enumerate(preds):
59             for gi, g in enumerate(gts):
60                 iou = interval_iou(p, g)
61                 # Require real overlap: at iou_threshold=0.0 a plain `iou >= threshold`
62                 # match completely disjoint intervals (IoU 0.0), fabricating perfect scores.
63                 if iou > 0.0 and iou >= iou_threshold:
64                     candidates.append((iou, pi, gi))
65                 # Highest IoU first; ties broken by indices for determinism.
66                 candidates.sort(key=lambda c: (-c[0], c[1], c[2]))
67
68         used_pred, used_gt = set(), set()
69         matches: List[Match] = []
70         for iou, pi, gi in candidates:
71             if pi in used_pred or gi in used_gt:
72                 continue
73             used_pred.add(pi)
74             used_gt.add(gi)
75             matches.append(Match(pred_index=pi, gt_index=gi, iou=iou))
76
77         unmatched_pred = [i for i in range(len(preds)) if i not in used_pred]
78         unmatched_gt = [i for i in range(len(gts)) if i not in used_gt]
79         return MatchResult(matches=matches,
80                             unmatched_pred_indices=unmatched_pred,
81                             unmatched_gt_indices=unmatched_gt)
82
83
84     @dataclass
85     class Score:
86         """Aggregate metrics for one (predictions vs ground-truth) evaluation."""
87         iou_threshold: float
88         num_pred: int
89         num_gt: int
90         true_positives: int
91         false_positives: int
92         false_negatives: int
93         precision: float
94         recall: float
95         f1: float
96         mean_iou: float # over matched pairs (0.0 if none)
97         mean_start_error: float # mean |pred_start - gt_start| over matches, seconds
98         mean_end_error: float # mean |pred_end - gt_end| over matches, seconds
99
100         def as_dict(self) -> dict:
101             return self.__dict__.copy()
102
103
104     def score(preds: List[Interval], gts: List[Interval],

```

```

105         iou_threshold: float = 0.5,
106         match_result: Optional[MatchResult] = None) -> Score:
107     """Compute precision/recall/F1, mean matched IoU, and boundary errors.
108
109     Pass a precomputed `match_result` to avoid re-matching (e.g. when also
110     extracting the failure lists); otherwise matching is done here.
111     """
112     mr = match_result if match_result is not None else match_intervals(preds, gts,
113     iou_threshold)
114
115     tp = len(mr.matches)
116     fp = len(mr.unmatched_pred_indices)
117     fn = len(mr.unmatched_gt_indices)
118
119     precision = tp / (tp + fp) if (tp + fp) > 0 else 0.0
120     recall = tp / (tp + fn) if (tp + fn) > 0 else 0.0
121     f1 = (2 * precision * recall / (precision + recall)) if (precision + recall) > 0
122     else 0.0
123
124     if mr.matches:
125         mean_iou = sum(m.iou for m in mr.matches) / tp
126         mean_start_err = sum(abs(preds[m.pred_index][0] - gts[m.gt_index][0]) for m in
127         mr.matches) / tp
128         mean_end_err = sum(abs(preds[m.pred_index][1] - gts[m.gt_index][1]) for m in
129         mr.matches) / tp
130     else:
131         mean_iou = mean_start_err = mean_end_err = 0.0
132
133     return Score(
134         iou_threshold=iou_threshold,
135         num_pred=len(preds), num_gt=len(gts),
136         true_positives=tp, false_positives=fp, false_negatives=fn,
137         precision=precision, recall=recall, f1=f1,
138         mean_iou=mean_iou, mean_start_error=mean_start_err, mean_end_error=mean_end_err,
139     )
140
141 def pr_curve(preds: List[Interval], gts: List[Interval],
142             thresholds: Optional[List[float]] = None) -> List[Score]:
143     """Score across a sweep of IoU thresholds (default 0.1..0.9 step 0.1).
144
145     Lets you see strict-vs-lenient behaviour rather than a single operating point.
146     """
147     if thresholds is None:
148         thresholds = [round(0.1 * i, 1) for i in range(1, 10)]
149     return [score(preds, gts, iou_threshold=t) for t in thresholds]

```

172. user (2026-06-13T20:41:55.355Z)

```

1  #!/usr/bin/env python
2  """Run the FULL test suite for the badminton-highlight-indexer.
3
4  A single, discoverable entry point so you (or a reviewer / CI) can verify the
5  whole suite before opening a PR:
6
7      python run_all_tests.py          # full suite, quiet
8      python run_all_tests.py -v       # verbose (passthrough to pytest)
9      python run_all_tests.py -k config # filter by keyword (passthrough)
10     python run_all_tests.py --cov=backend --cov-report=term-missing # coverage
11
12     Anything after the script name is passed straight through to pytest.
13
14     Notes
15     ----

```

```

16 - The suite is fully mocked (no GPU, no network, no WSL) and runs offline.
17 - The observability **reporting** tests skip cleanly if the optional ``obsreport``
18 package isn't installed (they use ``pytest.importorskip``). To exercise them too:
19     pip install -e ../sports-obsreport --no-build-isolation
20 - Exit code is pytest's own, so this is CI / pre-PR friendly (0 = all green).
21 """
22 from __future__ import annotations
23
24 import importlib.util
25 import subprocess
26 import sys
27
28
29 def _installed(module: str) -> bool:
30     try:
31         return importlib.util.find_spec(module) is not None
32     except (ImportError, ValueError):
33         return False
34
35
36 def main() -> int:
37     args = list(sys.argv[1:])
38     # Default to quiet unless the caller chose a verbosity.
39     if not any(a in ("-q", "-v", "-vv", "-vvv", "--quiet", "--verbose") for a in args):
40         args = ["-q", *args]
41     cmd = [sys.executable, "-m", "pytest", *args]
42
43     obs = "available" if _installed("obsreport") else "NOT installed ? reporting tests
will SKIP"
44     bar = "=" * 72
45     print(bar)
46     print(" Full test suite: " + " ".join(cmd))
47     print(f" obsreport (optional reporting dep): {obs}")
48     print(bar, flush=True)
49
50     rc = subprocess.run(cmd).returncode
51
52     print(bar)
53     print(" RESULT: " + ("PASS - all tests green" if rc == 0 else f"FAIL - pytest exit
{rc}"))
54     print(bar)
55     return rc
56
57
58 if __name__ == "__main__":
59     raise SystemExit(main())
60

```

173. user (2026-06-13T20:41:55.652Z)

```

1     name: CI
2
3     on:
4         push:
5             branches: [master]
6         pull_request:
7
8     permissions:
9         contents: read
10
11     concurrency:
12         group: ci-${{ github.ref }}
13         cancel-in-progress: true
14
15     jobs:

```

```

16 # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17 # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
18 # tracked .py, including scripts/ and training/, which the backend-only
19 # py_compile smoke sweep does not reach.
20 lint:
21   name: Ruff lint
22   runs-on: ubuntu-latest
23   steps:
24     - uses: actions/checkout@v6
25     - uses: actions/setup-python@v6
26       with:
27         python-version: "3.13"
28     - name: Install ruff
29       run: python -m pip install ruff
30     - name: Ruff check
31       run: ruff check . --output-format github
32
33 # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
34 # per-module quarantine list ? ratchet by removing entries as modules clean
35 # up. Typed core deps are installed so mypy sees their real signatures
36 # (without them ignore_missing_imports would Any-stub half the checks).
37 typecheck:
38   name: Mypy (lenient baseline)
39   runs-on: ubuntu-latest
40   steps:
41     - uses: actions/checkout@v6
42     - uses: actions/setup-python@v6
43       with:
44         python-version: "3.13"
45     - name: Install mypy + typed core deps
46       run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
genai numpy
47     - name: Mypy
48       run: mypy backend training
49
50 # Fast guard against the 7fbbeb0 class of breakage: a SyntaxError or
51 # module-level import error shipping while the suite stays green.
52 # Deliberately does NOT install cv2/torch/numpy ? the smoke tests must hold
53 # on machines without the GPU stack (test_import_smoke stubs what's absent).
54 import-smoke:
55   name: Syntax sweep + import smoke (no GPU stack)
56   runs-on: ubuntu-latest
57   steps:
58     - uses: actions/checkout@v6
59     - uses: actions/setup-python@v6
60       with:
61         python-version: "3.13"
62     - name: Install minimal deps
63       run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai
64     - name: Run import smoke tests
65       run: python -m pytest tests/test_import_smoke.py -v
66
67 full-suite:
68   name: Full test suite (offline, mocked)
69   runs-on: ubuntu-latest
70   steps:
71     - uses: actions/checkout@v6
72     - uses: actions/setup-python@v6
73       with:
74         python-version: "3.13"
75         cache: pip
76     - name: System libs for opencv-python
77       run: |
78         sudo apt-get update

```

```

79         sudo apt-get install -y libgl1 libglib2.0-0t64 || sudo apt-get install -y
libgl1 libglib2.0-0
80         - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
81         run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
82         - name: Install requirements
83         run: python -m pip install -r requirements.txt pytest pytest-cov httpx # httpx:
fastapi TestClient
84         # obsreport (private, soft dep) is absent here: reporting tests skip via
85         # pytest.importorskip ? run_all_tests.py prints which. The coverage floor
86         # is therefore calibrated against THIS lane's number (local runs with
87         # obsreport installed measure a little higher).
88         # Floor calibrated 2026-06-11 against this lane's measured 72% (local
89         # with obsreport: 73%). Ratchet it UP as coverage grows; never down
90         # without an owner decision.
91         - name: Run full suite (with coverage floor)
92         run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
93

```

174. user (2026-06-13T20:42:43.637Z)

```

1     # Lint-only lane (no `ruff format` enforcement ? a whole-repo reformat would
2     # bury blame history for zero behavior value).
3     #
4     # Selected: pyflakes + pycodestyle errors (ruff defaults) + pycodestyle
5     # warnings (W) + bugbear (B). Deferred with intent: I (import-sort) and UP
6     # (pyupgrade) ? both are large mechanical diffs; adopt via a dedicated PR.
7     target-version = "py38"
8
9     # scratch/ = throwaway debugging scripts that may hit real APIs; it is also
10    # excluded from pytest collection (see pytest.ini).
11    extend-exclude = ["scratch"]
12
13    [lint]
14    extend-select = ["W", "B"]
15
16    [lint.per-file-ignores]
17    # Entrypoint scripts deliberately run setup (load_dotenv, logging config,
18    # sys.path) before importing backend modules ? import order is load-bearing.
19    "backend/main.py" = ["E402"]
20    "run_phase3_eval.py" = ["E402"]
21    "run_process_and_stitch.py" = ["E402"]
22

```

175. assistant (2026-06-13T20:42:42.609Z)

numpy's local, CI auto-discovers `tests/` (full-suite has a **70% coverage floor**), and `metrics.py` is pure-Python. I'll keep the new modules pure-Python (stdlib only) so they also pass the no-numpy import-smoke lane. Let me read `calibration.py` (avoid name clashes), the ruff config (line length), and a sample test for conventions.

176. user (2026-06-13T20:42:43.584Z)

```

1     """Quantify rally-segmentation improvement ? and guard against overfitting.
2
3     This is the measurement layer for per-video / per-corpus calibration. It answers
4     two questions honestly:
5
6     1. """Did precision improve, by how much?""" ? `improvement_report` scores a
7       baseline config vs a tuned config on the same ground truth and reports the
8       before/after metrics and the % delta.
9
10    2. """Is that gain real or overfit?""" ? calibration must be measured on data the

```

```

11     tuning never saw, otherwise you're just memorising. Two protocols:
12     - `cross_validate` ? k-fold over one video's rallies: pick the config on the
13       train folds, score on the held-out fold. The **generalization_gap**
14       (train ? test) is the overfit alarm.
15     - `leave_one_video_out` ? the protocol that matters as more videos arrive:
16       pick the config on all *other* videos, score on the held-out video. Reports
17       per-video held-out scores + mean±std. This is the number to trust.
18
19     Everything is segmenter-agnostic: callers pass a `predict_fn(config) ->
20     List[Interval]` (e.g. WASB trajectory ? windows at those thresholds), so the same
21     harness works for WASB, yolo_hybrid, or anything else. Pure + unit-tested; no I/O.
22     """
23     from __future__ import annotations
24
25     import statistics as _st
26     from typing import Callable, Dict, List, Sequence, Tuple, Any
27
28     from backend.eval.metrics import score, Score, Interval
29
30     Config = Any                                # opaque to this module (e.g. a thresholds
dict)
31     PredictFn = Callable[[Config], List[Interval]]
32
33
34     def pct_improvement(before: float, after: float) -> float:
35         """Percentage change from `before` to `after` (e.g. 0.6 -> 0.75 == +25.0)."""
36         if before <= 0:
37             return float("inf") if after > 0 else 0.0
38         return (after - before) / before * 100.0
39
40
41     def evaluate(preds: List[Interval], gts: List[Interval], iou_threshold: float = 0.5) ->
Score:
42         """Precision / recall / F1 / mean-IoU of predicted vs ground-truth intervals."""
43         return score(preds, gts, iou_threshold)
44
45
46     def kfold_indices(n: int, k: int) -> List[Tuple[List[int], List[int]]]:
47         """Deterministic k-fold split of indices 0..n-1 into (train_idx, test_idx)."""
48         if n <= 0:
49             return []
50         k = max(2, min(k, n))                # need >=2 folds; cap at n
51         folds = [list(range(i, n, k)) for i in range(k)]    # round-robin ? balanced,
deterministic
52         splits = []
53         for f in range(k):
54             test = folds[f]
55             train = [i for i in range(n) if i not in set(test)]
56             if test and train:
57                 splits.append((train, test))
58         return splits
59
60
61     def select_best(predict_fn: PredictFn, configs: Sequence[Config], gts: List[Interval],
62                     metric: str = "f1", iou_threshold: float = 0.5) -> Tuple[Config, float]:
63         """Pick the config whose predictions maximise `metric` against `gts`."""
64         best_cfg, best_val = None, -1.0
65         for cfg in configs:
66             val = getattr(score(predict_fn(cfg), gts, iou_threshold), metric)
67             if val > best_val:
68                 best_cfg, best_val = cfg, val
69         return best_cfg, best_val
70
71
72     def cross_validate(predict_fn: PredictFn, configs: Sequence[Config], gts:

```

```

List[Interval],
73             k: int = 5, metric: str = "recall", iou_threshold: float = 0.5) ->
Dict[str, Any]:
74     """k-fold threshold calibration on ONE video's rallies (within-video overfit guard).
75
76     Per fold: choose the best config on the train rallies, score it on the held-out
77     rallies. A large `generalization_gap` (train >> test) means the thresholds are
78     overfitting that video's quirks.
79
80     Default metric is **recall** on purpose: predictions are global (whole-video
81     windows), so scoring them against a *fold subset* of GT would unfairly count the
82     other rallies' windows as false positives ? precision/F1 are not well-defined
83     per-fold. Recall ("did the tuned thresholds still detect the held-out rallies?")
84     is the honest within-video generalization signal. For precision/F1, use
85     `leave_one_video_out` (each video scored whole) or `improvement_report` on a
86     held-out video.
87     """
88     splits = kfold_indices(len(gts), k)
89     if not splits:
90         return {"error": "not enough rallies for cross-validation", "n": len(gts)}
91     train_scores, test_scores, picked = [], [], []
92     for train_idx, test_idx in splits:
93         gts_tr = [gts[i] for i in train_idx]
94         gts_te = [gts[i] for i in test_idx]
95         cfg, tr = select_best(predict_fn, configs, gts_tr, metric, iou_threshold)
96         te = getattr(score(predict_fn(cfg), gts_te, iou_threshold), metric)
97         train_scores.append(tr)
98         test_scores.append(te)
99         picked.append(cfg)
100     mean_test = _st.mean(test_scores)
101     mean_train = _st.mean(train_scores)
102     return {
103         "metric": metric,
104         "folds": len(splits),
105         "mean_test": mean_test,
106         "std_test": _st.pstdev(test_scores) if len(test_scores) > 1 else 0.0,
107         "mean_train": mean_train,
108         "generalization_gap": mean_train - mean_test,    # >~0.1 = likely overfitting
109         "picked_configs": picked,
110     }
111
112
113 def leave_one_video_out(videos: List[Dict[str, Any]], configs: Sequence[Config],
114                         metric: str = "f1", iou_threshold: float = 0.5) -> Dict[str,
Any]:
115     """Leave-one-video-out CV ? the honest metric as the corpus grows.
116
117     `videos`: list of {"name": str, "predict_fn": PredictFn, "gts": [Interval]}.
118     For each held-out video: pick the config that's best *pooled across the other
119     videos*, then score it on the held-out one. The mean held-out score is the
120     "precision on unseen videos" number; per-video spread shows stability.
121     """
122     if len(videos) < 2:
123         return {"error": "need >=2 videos for leave-one-video-out", "n": len(videos)}
124     per_video = []
125     for i, held in enumerate(videos):
126         others = [v for j, v in enumerate(videos) if j != i]
127
128         def pooled_score(cfg: Config, others: List[dict] = others) -> float:
129             vals = [getattr(score(v["predict_fn"](cfg), v["gts"], iou_threshold),
metric)
130                        for v in others]
131             return _st.mean(vals)
132
133     best_cfg, best_train = max(((c, pooled_score(c)) for c in configs), key=lambda

```

```

t: t[1])
134         held_score = getattr(score(held["predict_fn"])(best_cfg), held["gts"],
iou_threshold), metric)
135         per_video.append({"video": held.get("name", f"video_{i}"),
136                           "held_out": held_score, "train_pooled": best_train, "config":
best_cfg})
137     held_scores = [p["held_out"] for p in per_video]
138     train_scores = [p["train_pooled"] for p in per_video]
139     return {
140         "metric": metric,
141         "mean_held_out": _st.mean(held_scores),
142         "std_held_out": _st.pstdev(held_scores) if len(held_scores) > 1 else 0.0,
143         "generalization_gap": _st.mean(train_scores) - _st.mean(held_scores),
144         "per_video": per_video,
145     }
146
147
148     def improvement_report(predict_fn: PredictFn, baseline_config: Config, tuned_config:
Config,
149                           gts: List[Interval], iou_threshold: float = 0.5) -> Dict[str,
Any]:
150         """The "+X% precision" headline: baseline vs tuned config on the same GT.
151
152         NB: for an *honest* headline, pass GT the tuning never saw (a held-out video),
153         or read the cross-validated `mean_test` from `cross_validate` instead.
154         """
155         b = score(predict_fn(baseline_config), gts, iou_threshold)
156         t = score(predict_fn(tuned_config), gts, iou_threshold)
157         return {
158             "baseline": {"precision": b.precision, "recall": b.recall, "f1": b.f1,
"mean_iou": b.mean_iou},
159             "tuned": {"precision": t.precision, "recall": t.recall, "f1": t.f1, "mean_iou":
t.mean_iou},
160             "precision_delta_pts": t.precision - b.precision,
161             "precision_delta_pct": pct_improvement(b.precision, t.precision),
162             "f1_delta_pct": pct_improvement(b.f1, t.f1),
163         }
164

```

177. assistant (2026-06-13T20:44:16.586Z)

Conventions are clear (pure-Python like `metrics.py`, py38 target, ruff E/F/W/B with no line-length rule, mypy lenient, pytest auto-discovery, 70% coverage floor, and `calibration.py` already owns LOVO so I won't duplicate it). I'll add three isolated, stdlib-only modules + tests, keeping existing behavior untouched. Writing them now.

178. assistant (2026-06-13T20:44:45.208Z)

"""Segmentation-aware metrics (course-correction C4) ? make over-segmentation visible.

Temporal IoU (`metrics.score`) is nearly blind to *fragmentation*: a single rally chopped into three predicted pieces, or two rallies merged across a dead-time gap, can still post a moderate mean-IoU. The temporal-action-segmentation / -localization literature grades those failure modes with **F1@k** (segmental F1 at several overlap thresholds) and **mAP@tIoU** (ranked, confidence-aware average precision over a tIoU sweep). Reporting these alongside temporal-IoU is what lets a reconciler's split/merge/tighten win actually show up ? and stops an IoU-chasing rule from shredding rally structure invisibly. See `docs/ANALYZERS\_AND\_RECONCILERS.md` §13/C4.

This is **additive**: new functions over the existing `metrics` matcher; nothing here changes `score`/`compare`/LOVO behaviour. Pure (stdlib only), deterministic, no I/O ? so it runs in the GPU-free / numpy-free lanes too.

Conventions match `metrics`: an `Interval` is a `(start, end)` tuple of seconds. A `ScoredInterval` is `(interval, confidence)` ? the confidence is whatever the scorer

emits per predicted rally (e.g. the classifier probability), used only to *rank* for AP.

Note on the segmental *edit* score: the classic TAS edit/Levenshtein score operates on the ordered sequence of segment *labels* and is degenerate for a single action class (every rally carries the same label), so it is intentionally omitted ? F1@k + mAP@tIoU + the segment-count ratio are the over-segmentation-sensitive metrics for this binary setting.

```
from __future__ import annotations
```

```
from typing import Dict, List, Sequence, Tuple
```

```
from backend.eval.metrics import Interval, interval_iou, score
```

```
#: A predicted interval plus the confidence used to rank it for average precision.
```

```
ScoredInterval = Tuple[Interval, float]
```

```
#: Default overlap thresholds for F1@k ? the TAS convention F1@{10,25,50}.
```

```
DEFAULT_OVERLAPS: Tuple[float, ...] = (0.1, 0.25, 0.5)
```

```
#: Default tIoU sweep for mean average precision ? the action-detection convention.
```

```
DEFAULT_TIOUS: Tuple[float, ...] = (0.3, 0.4, 0.5, 0.6, 0.7)
```

```
def f1_at_overlaps(preds: List[Interval], gts: List[Interval],
                   overlaps: Sequence[float] = DEFAULT_OVERLAPS) -> Dict[float, float]:
    """Segmental F1 at each overlap threshold ({overlap: f1}) ? the TAS F1@k.
```

```
    Reuses the canonical greedy one-to-one matcher, so a fragmented prediction (3 pieces
    for 1 rally) scores extra false positives and a merged prediction misses a rally ?
    both drag F1 down where temporal-IoU alone would not flinch.
```

```
    """
```

```
    return {ov: score(preds, gts, iou_threshold=ov).f1 for ov in overlaps}
```

```
def average_precision(scored_preds: Sequence[ScoredInterval], gts: List[Interval],
                      iou_threshold: float = 0.5) -> float:
```

```
    """Average precision at a single tIoU (all-points / VOC-2010 interpolation).
```

```
    Predictions are ranked by descending confidence; each is a true positive iff it
    claims a still-unmatched ground-truth rally at IoU >= `iou_threshold` (highest
    available IoU wins, one GT per prediction). Returns AP in `[0, 1]`. `0.0` when
    there are no ground-truth rallies or no predictions reach the threshold.
```

```
    """
```

```
    n_gt = len(gts)
```

```
    if n_gt == 0 or not scored_preds:
```

```
        return 0.0
```

```
    order = sorted(range(len(scored_preds)),
```

```
                   key=lambda i: (-scored_preds[i][1], i)) # conf desc, stable by index
```

```
    matched_gt = set()
```

```
    precisions: List[float] = []
```

```
    recalls: List[float] = []
```

```
    cum_tp = 0
```

```
    cum_fp = 0
```

```
    for i in order:
```

```
        interval, _conf = scored_preds[i]
```

```
        best_iou, best_gi = 0.0, -1
```

```
        for gi, g in enumerate(gts):
```

```
            if gi in matched_gt:
```

```
                continue
```

```
            iou = interval_iou(interval, g)
```

```
            if iou >= iou_threshold and iou > best_iou:
```

```
                best_iou, best_gi = iou, gi
```

```
    if best_gi >= 0:
```

```
        matched_gt.add(best_gi)
```

```
        cum_tp += 1
```

```

        else:
            cum_fp += 1
            precisions.append(cum_tp / (cum_tp + cum_fp))
            recalls.append(cum_tp / n_gt)

# Monotone precision envelope (take the max precision to the right), then integrate
# over recall steps: AP = sum_k (recall_k - recall_{k-1}) * precision_env_k.
for i in range(len(precisions) - 2, -1, -1):
    precisions[i] = max(precisions[i], precisions[i + 1])
ap = 0.0
prev_recall = 0.0
for p, r in zip(precisions, recalls):
    if r > prev_recall:
        ap += p * (r - prev_recall)
        prev_recall = r
return ap

def mean_average_precision(scored_preds: Sequence[ScoredInterval], gts: List[Interval],
                           iou_thresholds: Sequence[float] = DEFAULT_TIOUS) -> float:
    """Mean of ``average_precision`` over a tIoU sweep (the standard mAP@tIoU)."""
    tious = list(iou_thresholds)
    if not tious:
        return 0.0
    return sum(average_precision(scored_preds, gts, t) for t in tious) / len(tious)

def segment_count_ratio(preds: List[Interval], gts: List[Interval]) -> float:
    """Predicted-to-GT segment-count ratio ? a crude fragmentation diagnostic.

    ``> 1`` hints over-segmentation (more predicted pieces than rallies), ``< 1`` hints
    merging/misses. Only meaningful with ground truth present; returns ``0.0`` if ``gts``
    is empty (use F1@k / mAP for the graded picture).
    """
    if not gts:
        return 0.0
    return len(preds) / len(gts)

def segmentation_report(scored_preds: Sequence[ScoredInterval], gts: List[Interval],
                        overlaps: Sequence[float] = DEFAULT_OVERLAPS,
                        iou_thresholds: Sequence[float] = DEFAULT_TIOUS) -> Dict[str, object]:
    """One-call bundle: F1@k + mAP@tIoU + per-tIoU AP + count ratio.

    ``scored_preds`` carry confidences for AP; F1@k uses the bare intervals. Returns a
    plain dict (JSON-friendly, string keys) so it can ride next to a ``compare`` result
    or land in the experiment leaderboard.
    """
    preds = [iv for iv, _conf in scored_preds]
    return {
        "f1_at_overlaps": {f"{ov:g}": v for ov, v in f1_at_overlaps(preds, gts,
overlaps).items()},
        "ap_at_tiou": {f"{t:g}": average_precision(scored_preds, gts, t) for t in
iou_thresholds},
        "mean_ap": mean_average_precision(scored_preds, gts, iou_thresholds),
        "segment_count_ratio": segment_count_ratio(preds, gts),
        "num_pred": len(preds),
        "num_gt": len(gts),
    }

```

179. user (2026-06-13T20:44:45.274Z)

File created successfully at: C:/Users/avidu/Projects/bhi-eval-wt/backend/eval/segmentation_metrics.py (file state is current in your context ? no need to Read it back)

180. user (2026-06-13T20:45:45.246Z)

File created successfully at: C:/Users/avidu/Projects/bhi-eval-wt/backend/eval/score_calibration.py (file state is current in your context ? no need to Read it back)

181. assistant (2026-06-13T20:45:16.163Z)

Module 1 done. Now the calibration module (C2) ? Platt + isotonic (PAVA), pure-Python.

182. assistant (2026-06-13T20:45:45.208Z)

"""Per-cue score calibration (course-correction C2) ? make ``confidence`` mean something.

A CNN softmax, an audio onset score, and an integer net-crossing count live on wildly different scales. Fed raw into the learned arbiter (the ADR-004 logistic regression), the most over-confident cue dominates and starves complementary cues. Calibrating each producer's score to a probability *before* it becomes a feature column makes the ``(value, confidence)`` semantics real and the arbiter's weights stable. See ``docs/ANALYZERS\_AND\_RECONCILERS.md`` §13/C2.

Two calibrators, both standard:

- **Platt scaling** ? fit ``p = sigmoid(a\*s + b)`` (parametric; good for monotone, roughly-sigmoidal cues; uses Platt's target smoothing to avoid 0/1 overfitting).
- **Isotonic regression** (pool-adjacent-violators) ? non-parametric monotone fit; more flexible, needs a bit more data.

Plus ``brier\_score`` and ``expected\_calibration\_error`` to *measure* calibration. All **additive**, pure (stdlib only), deterministic ? fits run in the numpy-free lanes too. Calibrators are tiny dataclasses (their params are JSON-friendly), so a fitted calibrator can be pinned alongside a variant the way the classifier is.

"""

```
from __future__ import annotations
```

```
import bisect
```

```
import math
```

```
from dataclasses import dataclass
```

```
from typing import List, Sequence, Tuple
```

```
_EPS = 1e-12
```

```
def _sigmoid(z: float) -> float:
    if z >= 0:
        return 1.0 / (1.0 + math.exp(-z))
    e = math.exp(z)
    return e / (1.0 + e)
```

```
def _clip01(p: float) -> float:
    return min(1.0 - _EPS, max(_EPS, p))
```

```
@dataclass(frozen=True)
```

```
class PlattCalibrator:
```

```
    """``p = sigmoid(a*score + b)``. Call it on a raw cue score to get a probability."""
```

```
    a: float
```

```
    b: float
```

```
    def __call__(self, s: float) -> float:
        return _clip01(_sigmoid(self.a * s + self.b))
```

```
@dataclass(frozen=True)
```

```
class IsotonicCalibrator:
```

```

"""Piecewise-constant monotone map from raw score to probability (PAVA fit).

`xs` are sorted training scores, `ys` the non-decreasing fitted probabilities;
a query returns the fitted value of the largest `x <= score` (clipped at the ends).
"""
xs: List[float]
ys: List[float]

def __call__(self, s: float) -> float:
    if not self.xs:
        return 0.5
    idx = bisect.bisect_right(self.xs, s) - 1
    if idx < 0:
        idx = 0
    return self.ys[idx]

def fit_platt(scores: Sequence[float], labels: Sequence[int],
              iters: int = 200, lr: float = 0.5) -> PlattCalibrator:
    """Fit Platt scaling by gradient descent on (smoothed) log-loss.

    Scores are standardised internally for a stable fit, then folded back so the returned
    `a`/`b` apply to the **raw** score. Platt target smoothing
    (`t+ = (N+ + 1)/(N+ + 2)`, `t- = 1/(N- + 2)`) keeps the fit off the 0/1 rails.
    Degenerate inputs (empty, zero-variance, single class) fall back to a sane constant map.
    """
    n = len(scores)
    if n == 0 or n != len(labels):
        return PlattCalibrator(a=0.0, b=0.0) # ? p = 0.5 everywhere
    mean = sum(scores) / n
    var = sum((s - mean) ** 2 for s in scores) / n
    sd = math.sqrt(var)
    n_pos = sum(1 for y in labels if y >= 0.5)
    n_neg = n - n_pos
    if sd <= _EPS or n_pos == 0 or n_neg == 0:
        # No usable score spread or one class: best constant probability = base rate.
        base = _clip01(n_pos / n)
        return PlattCalibrator(a=0.0, b=math.log(base / (1.0 - base)))
    t_pos = (n_pos + 1.0) / (n_pos + 2.0)
    t_neg = 1.0 / (n_neg + 2.0)
    z = [(s - mean) / sd for s in scores]
    targets = [t_pos if y >= 0.5 else t_neg for y in labels]

    a, b = 1.0, 0.0 # in standardised space
    for _ in range(max(1, iters)):
        ga = gb = 0.0
        for zi, t in zip(z, targets):
            p = _sigmoid(a * zi + b)
            d = p - t
            ga += d * zi
            gb += d
        a -= lr * ga / n
        b -= lr * gb / n
    # Fold standardisation back to raw score: z = (s-mean)/sd ? a*z+b = (a/sd)*s + (b -
    a*mean/sd).
    return PlattCalibrator(a=a / sd, b=b - a * mean / sd)

def _pava(ys: Sequence[float]) -> List[float]:
    """Pool-adjacent-violators: nearest non-decreasing fit to `ys` (in given order)."""
    blocks: List[List[float]] = [] # each [value, weight]
    for y in ys:
        blocks.append([float(y), 1.0])
        while len(blocks) >= 2 and blocks[-2][0] > blocks[-1][0]:
            v2, w2 = blocks.pop()

```

```

        v1, w1 = blocks.pop()
        blocks.append([(v1 * w1 + v2 * w2) / (w1 + w2), w1 + w2])
fitted: List[float] = []
for value, weight in blocks:
    fitted.extend([value] * int(round(weight)))
return fitted

def fit_isotonic(scores: Sequence[float], labels: Sequence[int]) -> IsotonicCalibrator:
    """Fit isotonic (monotone) calibration via PAVA on score-sorted labels."""
    n = len(scores)
    if n == 0 or n != len(labels):
        return IsotonicCalibrator(xs=[], ys=[])
    order = sorted(range(n), key=lambda i: (scores[i], i))
    xs = [float(scores[i]) for i in order]
    ys = [_clip01(v) for v in _pava([float(labels[i]) for i in order])]
    return IsotonicCalibrator(xs=xs, ys=ys)

def brier_score(probs: Sequence[float], labels: Sequence[int]) -> float:
    """Mean squared error of probabilities vs 0/1 labels (lower = better calibrated)."""
    n = len(probs)
    if n == 0 or n != len(labels):
        return 0.0
    return sum((p - y) ** 2 for p, y in zip(probs, labels)) / n

def expected_calibration_error(probs: Sequence[float], labels: Sequence[int],
                              n_bins: int = 10) -> float:
    """Expected calibration error: |accuracy - confidence| per equal-width bin, weighted."""
    n = len(probs)
    if n == 0 or n != len(labels):
        return 0.0
    bins = max(1, n_bins)
    bin_sum_p = [0.0] * bins
    bin_sum_y = [0.0] * bins
    bin_cnt = [0] * bins
    for p, y in zip(probs, labels):
        b = min(bins - 1, int(p * bins))
        bin_sum_p[b] += p
        bin_sum_y[b] += y
        bin_cnt[b] += 1
    ece = 0.0
    for i in range(bins):
        if bin_cnt[i]:
            conf = bin_sum_p[i] / bin_cnt[i]
            acc = bin_sum_y[i] / bin_cnt[i]
            ece += (bin_cnt[i] / n) * abs(acc - conf)
    return ece

def calibration_report(probs: Sequence[float], labels: Sequence[int],
                      n_bins: int = 10) -> Tuple[float, float]:
    """Convenience: ``(brier, ece)`` for a set of probabilities vs labels."""
    return brier_score(probs, labels), expected_calibration_error(probs, labels, n_bins)

```

183. user (2026-06-13T20:46:40.303Z)

File created successfully at: C:/Users/avidu/Projects/bhi-eval-wt/backend/eval/eval_stats.py
(file state is current in your context ? no need to Read it back)

184. assistant (2026-06-13T20:46:11.760Z)

Now the stats module (C1 bias-aware aggregation + C3 out-of-fold stacking) ? pure-Python.

185. assistant (2026-06-13T20:46:40.273Z)

```
"""Honest small-N evaluation statistics (course-corrections C1 + C3).
```

The ~6-match golden corpus + a regularised logistic arbiter is exactly the regime where a single leave-one-video-out *mean* lies: leave-one-out anti-correlates each fold's training and held-out label means at small N, and one lucky/unlucky fold swings the headline. The ``compare``/LOVO machinery already produces **per-video held-out scores** ? this module turns those into an honest summary:

- **C1 ? never a single mean.** ``mean\_with\_ci`` / ``bootstrap\_ci`` report a confidence interval; ``paired\_sign\_test`` and ``paired\_delta\_ci`` grade a B?A improvement by how many *videos* it wins on (the distribution-free test the repo already leans on: "sign test needs $n \geq 10$ for $p < 0.05$ "), so the promotion gate can't promote noise.
- **C3 ? honest stacking.** ``out\_of\_fold\_predictions`` / ``grouped\_folds`` give a learned producer's *out-of-fold* outputs so feeding them to the arbiter doesn't leak (group by match, never a held-out window from the same video).

See ``docs/ANALYZERS\_AND\_RECONCILERS.md`` §13/C1+C3. All **additive**, pure (stdlib only), deterministic (bootstrap takes an explicit ``seed``) ? nothing here changes existing ``compare``/``calibration.leave\_one\_video\_out`` behaviour; callers opt in.

```
"""
```

```
from __future__ import annotations
```

```
import math
```

```
import random
```

```
import statistics
```

```
from typing import Any, Callable, Dict, List, Optional, Sequence, Tuple
```

```
Interval = Tuple[float, float]
```

```
FitPredict = Callable[[List[int], List[int]], List[Any]]
```

```
def bootstrap_ci(values: Sequence[float], confidence: float = 0.95,
                 n_boot: int = 2000, seed: int = 0) -> Tuple[float, float]:
    """Percentile bootstrap CI for the mean of ``values``. Deterministic given ``seed``.

    ``n<=1`` returns a degenerate ``(v, v)`` / ``(0, 0)`` interval ? there is no spread to
    resample. With  $n=6$  (our corpus) this is the honest "how wide could the mean be?" band.
    """
    vals = [float(v) for v in values]
    n = len(vals)
    if n == 0:
        return (0.0, 0.0)
    if n == 1:
        return (vals[0], vals[0])
    rng = random.Random(seed)
    means: List[float] = []
    for _ in range(max(1, n_boot)):
        means.append(sum(vals[rng.randrange(n)] for _ in range(n)) / n)
    means.sort()
    alpha = (1.0 - confidence) / 2.0
    last = len(means) - 1
    lo = means[max(0, int(math.floor(alpha * last)))]
    hi = means[min(last, int(math.ceil((1.0 - alpha) * last)))]
    return (lo, hi)
```

```
def mean_with_ci(values: Sequence[float], confidence: float = 0.95,
                 n_boot: int = 2000, seed: int = 0) -> Dict[str, float]:
    """``{mean, std, n, ci_low, ci_high, confidence}`` ? report this, never a bare mean (C1)."""
    vals = [float(v) for v in values]
    n = len(vals)
    mean = statistics.mean(vals) if vals else 0.0
    std = statistics.pstdev(vals) if n > 1 else 0.0
```

```

    lo, hi = bootstrap_ci(vals, confidence, n_boot, seed)
    return {"mean": mean, "std": std, "n": float(n),
            "ci_low": lo, "ci_high": hi, "confidence": confidence}

def paired_sign_test(deltas: Sequence[float], eps: float = 0.0) -> Dict[str, float]:
    """Two-sided exact sign test over per-fold deltas (B ? A).

    ``wins`` = folds where B beat A by more than ``eps``; ``losses`` the reverse; ties
    excluded. ``p_value`` is the exact two-sided binomial tail under p=0.5 ? the
    distribution-free "did B win on enough *videos*?" check.
    """
    wins = sum(1 for d in deltas if d > eps)
    losses = sum(1 for d in deltas if d < -eps)
    ties = len(deltas) - wins - losses
    n_eff = wins + losses
    if n_eff == 0:
        p = 1.0
    else:
        k = min(wins, losses)
        tail = sum(math.comb(n_eff, j) for j in range(k + 1)) * (0.5 ** n_eff)
        p = min(1.0, 2.0 * tail)
    return {"wins": float(wins), "losses": float(losses), "ties": float(ties),
            "n_effective": float(n_eff), "p_value": p}

def paired_delta_ci(a_values: Sequence[float], b_values: Sequence[float],
                    confidence: float = 0.95, n_boot: int = 2000, seed: int = 0,
                    eps: float = 0.0) -> Dict[str, Any]:
    """Paired (by fold) B ? A summary: mean delta + bootstrap CI + the sign test (C1).

    ``a_values``/``b_values`` are per-fold held-out scores for variants A and B, aligned by
    fold (same video order). A promotion is honest when the delta CI clears 0 *and* the sign
    test agrees ? not when a single mean ticked up.
    """
    n = min(len(a_values), len(b_values))
    deltas = [float(b_values[i]) - float(a_values[i]) for i in range(n)]
    summary = mean_with_ci(deltas, confidence, n_boot, seed)
    return {
        "mean_delta": summary["mean"],
        "ci_low": summary["ci_low"],
        "ci_high": summary["ci_high"],
        "n": float(n),
        "ci_excludes_zero": (summary["ci_low"] > 0.0) or (summary["ci_high"] < 0.0),
        "sign_test": paired_sign_test(deltas, eps),
    }

def grouped_folds(groups: Sequence[Any]) -> List[Tuple[List[int], List[int]]]:
    """Leave-one-group-out splits ``(train_idx, test_idx)`` ? group by match, the honest unit.

    The field's name for "LOVO by match" is leave-one-group-out (LOGO). Group order is
    first-seen (deterministic). Groups with no complement (a single group) yield no fold.
    """
    seen: List[Any] = []
    for g in groups:
        if g not in seen:
            seen.append(g)
    n = len(groups)
    folds: List[Tuple[List[int], List[int]]] = []
    for g in seen:
        test = [i for i in range(n) if groups[i] == g]
        train = [i for i in range(n) if groups[i] != g]
        if test and train:
            folds.append((train, test))

```

```
return folds
```

```
def out_of_fold_predictions(groups: Sequence[Any], fit_predict: FitPredict) ->
List[Optional[Any]]:
    """Out-of-fold predictions for honest stacking (C3).

    For each leave-one-group-out fold, ``fit_predict(train_idx, test_idx)`` trains on the
    other groups and returns predictions for ``test_idx`` (in that order); results are
    scattered back to a list aligned to the input items. So a *learned* producer's outputs,
    when used as arbiter features, never include in-sample predictions ? the leakage path
    that inflates measured lift on a tiny corpus. Items in a singleton group stay ``None``
    (needs >= 2 groups to predict).
    """
    n = len(groups)
    oof: List[Optional[Any]] = [None] * n
    for train_idx, test_idx in grouped_folds(groups):
        preds = fit_predict(train_idx, test_idx)
        for pos, ti in enumerate(test_idx):
            if pos < len(preds):
                oof[ti] = preds[pos]
    return oof
```

186. assistant (2026-06-13T20:46:47.344Z)

Now the tests. Let me match the repo's existing test style.

187. user (2026-06-13T20:46:49.500Z)

```
test_calibrate_local.py
test_calibrate_wasb.py
test_calibration.py
test_distill_local.py
test_eval_metrics.py
test_experiment.py
--- sample test header (test_metrics.py if present) ---
"""Tests for the experiment harness (backend.eval.experiment).

Pure/offline ? synthetic candidate rows, no GPU/WSL. The load-bearing test is the
**no-regression invariant** (``test_disabled_cues_byte_identical_to_control``): a variant
that merely carries a disabled cue must predict exactly what control predicts. That is the
guarantee that lets new cues be merged freely (``EXPERIMENT_HARNESS.md`` §6).
"""
import json

import pytest

from backend.eval.classifier import LogisticRegression
from backend.eval.experiment import (
    CONTROL,
    ExperimentError,
    Variant,
    VideoCandidates,
    compare,
    compare_unlabeled,
    evaluate_variant,
    load_video_candidates,
    predict,
    record_experiment,
)
from backend.infrastructure.database import Database
```

A video whose third candidate is a held-shuttle false fire (net_crossings=0): it does

not overlap any golden rally, so it is a negative by construction.

```
def _held_shuttle_video(name="v", gts=None):
```

188. assistant (2026-06-13T20:47:21.972Z)

Now the three test files, matching the repo's pytest style.

189. assistant (2026-06-13T20:47:34.605Z)

```
"""Tests for backend.eval.segmentation_metrics (course-correction C4).

Pure/offline. The point of these metrics is to *see over-segmentation* that temporal-IoU
hides, so the tests pin: perfect = 1.0, fragmentation and merging drop F1@k, and AP
respects confidence ranking.
"""

import pytest

from backend.eval.segmentation_metrics import (
    average_precision,
    f1_at_overlaps,
    mean_average_precision,
    segment_count_ratio,
    segmentation_report,
)

GTS = [(0.0, 10.0), (20.0, 30.0), (40.0, 50.0)]

def test_f1_perfect_is_one_at_every_overlap():
    f1 = f1_at_overlaps(GTS, GTS, overlaps=(0.1, 0.25, 0.5))
    assert f1 == {0.1: 1.0, 0.25: 1.0, 0.5: 1.0}

def test_fragmentation_drops_f1():
    # One rally split into two sub-segments ? an extra false positive at strict overlap.
    frag = [(0.0, 4.0), (5.0, 10.0), (20.0, 30.0), (40.0, 50.0)]
    assert f1_at_overlaps(frag, GTS)[0.5] < 1.0

def test_over_merge_drops_f1():
    # Two rallies merged across the dead-time gap ? low IoU with either, both missed.
    merged = [(0.0, 30.0), (40.0, 50.0)]
    assert f1_at_overlaps(merged, GTS)[0.5] < 1.0

def test_average_precision_perfect_ranking_is_one():
    scored = [((0.0, 10.0), 0.9), ((20.0, 30.0), 0.8), ((40.0, 50.0), 0.7)]
    assert average_precision(scored, GTS, iou_threshold=0.5) == pytest.approx(1.0)

def test_average_precision_penalises_high_confidence_false_positive():
    good = [((0.0, 10.0), 0.9), ((20.0, 30.0), 0.8), ((40.0, 50.0), 0.7)]
    with_fp = [((100.0, 110.0), 0.95)] + good          # an FP ranked first
    assert average_precision(with_fp, GTS, 0.5) < average_precision(good, GTS, 0.5)

def test_average_precision_degenerate_inputs():
    assert average_precision([], GTS) == 0.0          # no predictions
    assert average_precision([(0.0, 10.0), 0.9], []) == 0.0  # no ground truth

def test_mean_average_precision_averages_the_sweep():
    scored = [((0.0, 10.0), 0.9), ((20.0, 30.0), 0.8), ((40.0, 50.0), 0.7)]
    assert mean_average_precision(scored, GTS, iou_thresholds=(0.3, 0.5, 0.7)) ==
```

```

pytest.approx(1.0)
    assert mean_average_precision(scored, GTS, iou_thresholds=()) == 0.0

def test_segment_count_ratio_flags_fragmentation():
    assert segment_count_ratio([(0, 1), (2, 3), (4, 5), (6, 7), (8, 9)], GTS) == pytest.approx(5
/ 3)
    assert segment_count_ratio([(0, 1)], []) == 0.0    # undefined without GT

def test_segmentation_report_shape():
    scored = [((0.0, 10.0), 0.9), ((20.0, 30.0), 0.8), ((40.0, 50.0), 0.7)]
    rep = segmentation_report(scored, GTS)
    assert rep["num_pred"] == 3 and rep["num_gt"] == 3
    assert rep["mean_ap"] == pytest.approx(1.0)
    assert set(rep["f1_at_overlaps"]) == {"0.1", "0.25", "0.5"}
    assert rep["f1_at_overlaps"]["0.5"] == pytest.approx(1.0)

```

190. user (2026-06-13T20:47:34.626Z)

File created successfully at: C:/Users/avidu/Projects/bhi-eval-wt/tests/test_segmentation_metrics.py (file state is current in your context ? no need to Read it back)

191. assistant (2026-06-13T20:47:49.139Z)

```

"""Tests for backend.eval.score_calibration (course-correction C2).

```

Pure/offline. Pins: calibrators are monotone and output probabilities in [0,1]; Brier/ECE behave; degenerate inputs fall back sanely.

```

"""

```

```

import pytest

```

```

from backend.eval.score_calibration import (
    brier_score,
    calibration_report,
    expected_calibration_error,
    fit_isotonic,
    fit_platt,
)

```

```

def test_platt_is_monotone_and_separates():
    scores = [-3.0, -2.0, -1.0, 1.0, 2.0, 3.0]
    labels = [0, 0, 0, 1, 1, 1]
    cal = fit_platt(scores, labels)
    assert cal(-5.0) < cal(5.0)    # positive correlation ? increasing
    assert 0.0 < cal(-5.0) < 0.5 < cal(5.0) < 1.0
    assert all(0.0 <= cal(s) <= 1.0 for s in scores)

```

```

def test_platt_degenerate_inputs():
    assert fit_platt([], [])(0.0) == pytest.approx(0.5)    # empty ? 0.5
    assert fit_platt([1.0, 1.0], [1, 1])(1.0) > 0.5    # one class ? base rate > 0.5
    assert fit_platt([5.0, 5.0], [1, 0])(5.0) == pytest.approx(0.5)    # no spread ? base rate 0.5

```

```

def test_isotonic_is_monotone_and_pools_violations():
    cal = fit_isotonic([1.0, 2.0, 3.0, 4.0], [0, 0, 1, 1])
    assert cal(1.0) < cal(4.0)
    assert cal(0.0) == cal(1.0)    # below range clips to first
    assert cal(99.0) == cal(4.0)    # above range clips to last
    # A violation (1,0,1) must be pooled into a non-decreasing fit.
    pooled = fit_isotonic([1.0, 2.0, 3.0], [1, 0, 1])
    assert pooled(1.0) <= pooled(2.0) <= pooled(3.0)

```

```

def test_isotonic_empty():
    assert fit_isotonic([], [])(0.0) == pytest.approx(0.5)

def test_brier_score_known_values():
    assert brier_score([1.0, 0.0], [1, 0]) == pytest.approx(0.0)
    assert brier_score([0.5, 0.5], [1, 0]) == pytest.approx(0.25)

def test_ece_zero_when_calibrated_positive_when_not():
    # Perfectly calibrated: prob 1.0 with label 1, prob 0.0 with label 0.
    assert expected_calibration_error([1.0, 1.0, 0.0, 0.0], [1, 1, 0, 0]) == pytest.approx(0.0)
    # Confident but wrong ? large gap.
    assert expected_calibration_error([0.9, 0.9], [0, 0]) > 0.5

def test_calibration_report_returns_pair():
    brier, ece = calibration_report([1.0, 0.0], [1, 0])
    assert brier == pytest.approx(0.0) and ece == pytest.approx(0.0)

```

192. user (2026-06-13T20:47:49.188Z)

File created successfully at: C:/Users/avidu/Projects/bhi-eval-wt/tests/test_score_calibration.py (file state is current in your context ? no need to Read it back)

193. assistant (2026-06-13T20:48:06.981Z)

```

"""Tests for backend.eval.eval_stats (course-corrections C1 + C3).

Pure/offline + deterministic (the bootstrap takes a fixed seed). Pins: CIs bracket the
mean and degrade sanely at tiny n; the sign test's exact p-value; leave-one-group-out
folds; and that out-of-fold predictions are genuinely out of fold (no leakage).
"""

import pytest

from backend.eval.eval_stats import (
    bootstrap_ci,
    grouped_folds,
    mean_with_ci,
    out_of_fold_predictions,
    paired_delta_ci,
    paired_sign_test,
)

def test_bootstrap_ci_is_deterministic_and_brackets_mean():
    vals = [0.5, 0.6, 0.7, 0.55, 0.65, 0.6]
    lo1, hi1 = bootstrap_ci(vals, seed=0)
    lo2, hi2 = bootstrap_ci(vals, seed=0)
    assert (lo1, hi1) == (lo2, hi2) # same seed ? identical
    assert lo1 <= sum(vals) / len(vals) <= hi1

def test_bootstrap_ci_degenerate():
    assert bootstrap_ci([]) == (0.0, 0.0)
    assert bootstrap_ci([0.42]) == (0.42, 0.42)
    assert bootstrap_ci([0.5, 0.5, 0.5]) == (0.5, 0.5) # no spread ? point interval

def test_mean_with_ci_shape():
    out = mean_with_ci([0.4, 0.6, 0.5], seed=1)
    assert out["mean"] == pytest.approx(0.5)
    assert out["n"] == 3.0
    assert out["ci_low"] <= out["mean"] <= out["ci_high"]

```

```

def test_sign_test_all_wins_exact_p():
    res = paired_sign_test([0.1, 0.2, 0.05, 0.3, 0.15, 0.2]) # 6 wins
    assert res["wins"] == 6.0 and res["losses"] == 0.0
    assert res["p_value"] == pytest.approx(2 * 0.5 ** 6) # 0.03125

def test_sign_test_ties_and_mixed():
    res = paired_sign_test([0.0, 0.0, 0.0])
    assert res["ties"] == 3.0 and res["n_effective"] == 0.0 and res["p_value"] == 1.0
    mixed = paired_sign_test([0.1, -0.1, 0.2])
    assert mixed["wins"] == 2.0 and mixed["losses"] == 1.0

def test_paired_delta_ci_detects_consistent_improvement():
    a = [0.50, 0.55, 0.60, 0.52, 0.58, 0.54]
    b = [0.60, 0.66, 0.69, 0.61, 0.70, 0.63] # B beats A every fold
    out = paired_delta_ci(a, b, seed=0)
    assert out["mean_delta"] > 0.0
    assert out["sign_test"]["wins"] == 6.0
    assert out["ci_excludes_zero"] is True

def test_grouped_folds_leave_one_group_out():
    groups = ["m1", "m1", "m2", "m2", "m3"]
    folds = grouped_folds(groups)
    assert len(folds) == 3
    train0, test0 = folds[0] # first-seen group m1 held out
    assert test0 == [0, 1] and train0 == [2, 3, 4]
    assert grouped_folds(["only"]) == [] # single group ? no fold

def test_out_of_fold_predictions_are_out_of_fold():
    groups = ["a", "a", "b", "b"]
    y = [0, 0, 1, 1]

    def fit_predict(train_idx, test_idx):
        # "Model" = mean label of the training groups; returned for each test item.
        m = sum(y[i] for i in train_idx) / len(train_idx)
        return [m] * len(test_idx)

    oof = out_of_fold_predictions(groups, fit_predict)
    # Group a's OOF uses only group b (mean 1.0) and vice-versa ? never its own labels.
    assert oof == [1.0, 1.0, 0.0, 0.0]

```

194. user (2026-06-13T20:48:07.009Z)

File created successfully at: C:/Users/avidu/Projects/bhi-eval-wt/tests/test_eval_stats.py (file state is current in your context ? no need to Read it back)

195. assistant (2026-06-13T20:48:13.152Z)

Now run the new tests locally to verify correctness before pushing.

196. user (2026-06-13T20:48:16.094Z)

```

..... [100%]
24 passed in 0.32s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

197. assistant (2026-06-13T20:48:24.106Z)

All 24 pass. Now the lint + type gates (CI runs ruff and `mypy backend training`). Let me run them on the new files.

198. user (2026-06-13T20:48:30.790Z)

```
=== ruff ===
All checks passed!
=== mypy (new modules) ===
Success: no issues found in 3 source files
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

199. assistant (2026-06-13T20:48:48.107Z)

Ruff + mypy clean. Let me run mypy exactly as CI does (`mypy backend training` with the repo's `mypy.ini`) to be certain no baseline regression, then I'll commit. (The full pytest suite needs torch/cv2 which aren't installed locally, but my change is purely additive ? 3 new modules + 3 new tests, no existing file touched ? so it can't break existing tests, and the new modules are well-covered so coverage rises.)

200. user (2026-06-13T20:50:09.931Z)

```
backend\annotations\base.py:85: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\sports\__init__.py:7: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:36: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 88 source files
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

201. assistant (2026-06-13T20:50:20.456Z)

mypy clean (88 files, the notes are pre-existing baseline). Now the "updated docs" part ? add an entry to the living quality log. Let me read its top to place the entry correctly.

202. user (2026-06-13T20:50:21.292Z)

```
40
41     ## Current track (live) ? newest first
42
43     ### Multi-signal fusion P2: the FusionSegmenter *(in progress, 2026-06-14)*
44     **Hypothesis.** Adjudicating shuttle-seeded windows with the net-crossing **Gate A**
(`min_crossings?2`)
45     and a person-occupancy **Gate B** raises **precision** on the held-shuttle / warm-up /
one-sided-feed /
46     spectator false positives ? *without* hurting recall, because the shuttle stays the
primary,
47     camera-invariant seed.
48
49     **Scope ? and an honesty boundary from the archive (read this).** The
50     [`2026-06-multivideo-lovo`](archives/quality-iterations/2026-06-multivideo-
lovo/README.md) iteration
51     already smoke-tested **court-region detection as a fix for the *multi-court contrast*
confound** and
52     found it **low-headroom**: only 12?14% of dead-time shuttle detections fall *outside*
the foreground
53     court, so a perfect court gate lifts contrast to only ~1.6?2.2× (below the ?3.6× "works"
threshold) ? the
54     multi-court problem is **temporal**, and the **learned model**, not a court mask, is its
answer. So Gate B
55     here is **NOT** sold as the multi-court fix. It targets a **different, previously-
untested** failure mode
56     (held-shuttle/warm-up/one-sided/spectator FPS). Two distinct levers; don't conflate
them.
57
```

```

58      **Data support for Gate A.** The same archive shows real rallies on the multi-court
matches contain
59      **?2 net-crossings 97% of the time (median 5)** while a single service feed crosses
~once ? `min_crossings?2`
60      is a principled, data-backed discriminator (it's the held-shuttle owner-bug killer).
61
62      **Shipping shape.** A NEW `fusion` segmenter, **default OFF** (`default_segmenter`
unchanged); Gate A
63      wired into the live trajectory path behind `indexing.min_crossings` (default 0 = OFF ?
wasb/tracknet
64      byte-identical today); the court mask is an **optional** config polygon (omit ? Gate B
is
65      player-count-only; supply ? masked). Person features (`players_in_court,
two_sided_motion,
66      mean_player_motion, in_court_ratio, shuttle_player_colocation`) persisted into candidate
`stats` for P3's
67      learned fusion. Reuses `PersonDetector` (P1) + `rally_gate` (already existed) ? no
reinvention.
68
69      **Status (2026-06-14): SHIPPED, default-OFF.** `FusionSegmenter`
(`backend/pipeline/segmenters/fusion_hybrid.py`,
70      registry key `fusion_hybrid`) + the pure feature math
(`backend/pipeline/detectors/fusion_features.py`) + two
71      identity hooks on the trajectory engine (`_enrich_candidate_stats`,
`_select_for_handoff`) + `PersonDetector`.
72      detections_per_frame` + the `FusionConfig` block. Built design-first (a
`fusion-p2-design` fan-out workflow), then
73      **adversarially reviewed by a 31-agent workflow (25 findings ? 19 confirmed)**: the
headline confirmation is
74      **NO-REGRESSION verified** ? wasb/tracknet twins persist byte-identical stats and hand
off all windows (pinned by

```

203. assistant (2026-06-13T20:50:27.545Z)

I'll add a new entry at the top of the live track (newest-first) documenting the C1?C4 eval-honesty helpers.

204. user (2026-06-13T20:50:42.043Z)

The file C:/Users/avidu/Projects/bhi-eval-wt/docs/QUALITY_ITERATIONS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

205. assistant (2026-06-13T20:50:50.449Z)

All gates pass locally. Let me commit, push, and open the PR.

206. user (2026-06-13T20:51:03.764Z)

```

=== status ===
M docs/QUALITY_ITERATIONS.md
?? backend/eval/eval_stats.py
?? backend/eval/score_calibration.py
?? backend/eval/segmentation_metrics.py
?? tests/test_eval_stats.py
?? tests/test_score_calibration.py
?? tests/test_segmentation_metrics.py
warning: in the working copy of 'backend/eval/eval_stats.py', LF will be replaced by CRLF the
next time Git touches it
warning: in the working copy of 'backend/eval/score_calibration.py', LF will be replaced by CRLF
the next time Git touches it
warning: in the working copy of 'backend/eval/segmentation_metrics.py', LF will be replaced by
CRLF the next time Git touches it
warning: in the working copy of 'tests/test_eval_stats.py', LF will be replaced by CRLF the next
time Git touches it

```

```
warning: in the working copy of 'tests/test_score_calibration.py', LF will be replaced by CRLF
the next time Git touches it
warning: in the working copy of 'tests/test_segmentation_metrics.py', LF will be replaced by
CRLF the next time Git touches it
[feat/eval-honesty-corrections 181ab3a] feat(eval): measurement-honesty helpers (course
corrections C1-C4, additive/default-OFF)
 7 files changed, 696 insertions(+)
 create mode 100644 backend/eval/eval_stats.py
 create mode 100644 backend/eval/score_calibration.py
 create mode 100644 backend/eval/segmentation_metrics.py
 create mode 100644 tests/test_eval_stats.py
 create mode 100644 tests/test_score_calibration.py
 create mode 100644 tests/test_segmentation_metrics.py
=== push ===
branch 'feat/eval-honesty-corrections' set up to track 'origin/feat/eval-honesty-corrections'.
To https://github.com/avidullu/badminton-highlight-indexer.git
 * [new branch]      feat/eval-honesty-corrections -> feat/eval-honesty-corrections
```

207. assistant (2026-06-13T20:51:07.534Z)

Pushed (7 files, 696 insertions). Opening the PR.

208. user (2026-06-13T20:51:23.260Z)

<https://github.com/avidullu/badminton-highlight-indexer/pull/124>

209. assistant (2026-06-13T20:51:32.098Z)

PR [#124](<https://github.com/avidullu/badminton-highlight-indexer/pull/124>) opened. Now I'll wait for CI and merge on green (per your choice). Watching the checks.

210. user (2026-06-13T20:53:17.334Z)

Full test suite (offline, mocked) pending 0 <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698339>
Refreshing checks status every 20 seconds. Press Ctrl+C to quit.

```
Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698349
Ruff lint pass 6s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698345
Syntax sweep + import smoke (no GPU stack) pass 19s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698343
Full test suite (offline, mocked) pending 0 https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698339
Full test suite (offline, mocked) pass 1m32s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698339
Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698349
Ruff lint pass 6s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698345
Syntax sweep + import smoke (no GPU stack) pass 19s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698343
Full test suite (offline, mocked) pass 1m32s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698339
Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698349
Ruff lint pass 6s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698345
Syntax sweep + import smoke (no GPU stack) pass 19s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27478758795/job/81222698343
---exit: 0
```

211. assistant (2026-06-13T20:53:24.339Z)

All four CI lanes green (full suite 1m32s ? coverage floor held). Merging #124.

212. user (2026-06-13T20:53:33.352Z)

```
{"mergeStateStatus":"CLEAN","mergeable":"MERGEABLE"}
failed to delete local branch feat/eval-honesty-corrections: failed to run git: error: cannot
delete branch 'feat/eval-honesty-corrections' used by worktree at 'C:/Users/avidu/Projects/bhi-
eval-wt'

=== confirm ===
{"mergedAt":"2026-06-13T20:53:29Z","state":"MERGED"}
2c8b1ad feat(eval): measurement-honesty helpers (course corrections C1-C4, additive/default-OFF)
(#124)
f4d7008 feat(eval): fusion_golden progress logging + resumability (observability) (#123)
23672f2 docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework (warm-up +
shuttle-state) (#119)
```

213. assistant (2026-06-13T20:53:38.561Z)

Merged to master (`2c8b1ad`). Now clean up the worktree (which still holds the local branch).

214. user (2026-06-13T20:53:41.393Z)

```
worktree removed
Deleted branch feat/eval-honesty-corrections (was 181ab3a).
=== worktrees ===
C:/Users/avidu/Projects/badminton-highlight-indexer f4d7008 [master]
=== main checkout untouched? ===
master
    (only untracked carryovers)
```

215. user (2026-06-13T20:53:58.478Z)

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (RUN.LOG READABILITY for manual QA ? **PR #120 MERGED ?
master `6076eef`**; owner asked ? self-merged
13     on verified-green CI per [[working-agreement-merging]] (squash; remote+local branch
deleted); was branch `feat/runlog-readability`
14     off `35a50f3`; CODE, minimal/config-gated). Owner's manual-QA loop = run the tool on a
video ?
```

216. assistant (2026-06-13T20:53:57.730Z)

Everything's clean: PR #124 merged to `master` (`2c8b1ad`), worktree removed, local branch deleted, main checkout untouched. Let me record this checkpoint in memory (two PRs landed this session).

217. user (2026-06-13T20:54:21.976Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

218. assistant (2026-06-13T20:54:45.256Z)

Both done and on `master`. Here's the wrap-up.

Merged (self-merge on green CI, as you chose)

PR [#119](https://github.com/avidullu/badminton-highlight-indexer/pull/119) ? `master 23672f2` ? the framework doc + the **literature review** (§13): honest verdict, terminology realignment, course corrections, and verified citations.

PR [#124](https://github.com/avidullu/badminton-highlight-indexer/pull/124) ? `master 2c8b1ad` ? the **C1?C4 measurement-honesty helpers**, forked into its own branch as you asked, implemented **additive / default-OFF** (no existing `compare`/LOVO`metrics` behaviour touched):

```
Fix	Module	Adds
C1 ? never a bare mean	`backend/eval/eval_stats.py`	`mean_with_ci`/`bootstrap_ci`
(seeded), `paired_sign_test`, `paired_delta_ci`		
C2 ? calibrate cue confidence	`backend/eval/score_calibration.py`	Platt + isotonic (PAVA),
`brier_score`/`ECE`		
C3 ? honest stacking	`backend/eval/eval_stats.py`	`out_of_fold_predictions` /
`grouped_folds` (LOGO by match)		
C4 ? see over-segmentation	`backend/eval/segmentation_metrics.py`	F1@{10,25,50}, mAP@tIoU,
`segment_count_ratio` |
```

Docs + CI

- **Docs:** living-log entry in `docs/QUALITY\_ITERATIONS.md` + the §13 review already in #119.
- **CI:** all four lanes green on #124 ? and the **full-suite** lane now runs the 24 new tests (coverage floor held), so CI covers the new code. Pure-stdlib, so they also pass the no-numpy import-smoke lane. (I read "updated CI" as **CI** exercises the new work ? it does. If you meant a **new dedicated lane** ? e.g. a docs link-checker, or a separate eval-metrics job ? say the word and I'll add it.)

How I avoided the chaos

Both PRs were built in **isolated git worktrees off master**, never flipping the shared checkout ? siblings merged #117/#118/#120/#123 concurrently throughout, and #116 stayed closed (its content had already landed via #117). The main checkout was left untouched.

The one deliberate boundary (owner-gated follow-on)

I did **not** wire C1?C4 into the **default** harness output ? making `compare()` / `run\_regression.py` report CIs + F1@k instead of a bare mean **changes the numbers we report**, so it warrants its own PR and your eyes. That, plus the Tier-2/3 corrections (global reconciler decode, cue-correlation modeling, temporal smoothing, external ShuttleSet eval), is queued in §13. Want me to take the default-harness wiring next?

219. user (2026-06-14T03:02:28.676Z)

for the default harness, can we run a A/B experiment and see the effects before or a single video? Any small video which is golden would be fine

220. assistant (2026-06-14T03:03:12.876Z)

Good question ? and the whole point of the "persist detection once, re-score offline" design is that an A/B like this **should** be runnable here without a GPU, **if** a golden video's persisted candidates (or WASB trajectory) + golden labels are present locally. Let me first check what golden data actually exists in this checkout (I won't fabricate numbers ? if the data isn't here, this becomes a ready-to-run script for your machine).

221. user (2026-06-14T03:03:20.510Z)

```
=== golden manifest ===
[
  {
```

```

    "name": "adarsh_avi_singles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/adarsh_avi_full_traj.csv",
    "fps": 59.94, "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden Labelled/Adarsh and
Avi.rallies.csv"
  },
  {
    "name": "kushagra_singles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/kushagra_traj.csv",
    "fps": 59.94, "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/KushagraYashAviFirstVideo.rallies.csv"
  },
  {
    "name": "largetest_doubles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/largetestvideo_traj.csv",
    "fps": 29.97, "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/LargeTestVideo.rallies.csv"
  },
  {
    "name": "gbaaddy",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/gbaaddy_traj.csv",
    "fps": 59.94, "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/gBaaddy20May2026.rallies.csv"
  },
  {
    "name": "testlarge_short",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/testlarge_traj.csv",
    "fps": 29.97, "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/TestLargeVideo.rallies.csv"
  },
  {
    "name": "mahadevpura_singles",
    "trajectory": "C:/Users/avidu/Projects/Annotation
Setup/Collect/Trajectories/mahadevpura_traj.csv",
    "fps": 29.97, "frame_width": 1920,
    "labels": "C:/Users/avidu/Projects/Annotation Setup/Collect/Golden
Labelled/MahadevpuraSingles.rallies.csv"
  }
]

```

=== DBs present ===

```

./..mypy_cache/3.13/cache.0.db
./..mypy_cache/3.13/cache.1.db
./..mypy_cache/3.13/cache.10.db
./..mypy_cache/3.13/cache.11.db
./..mypy_cache/3.13/cache.12.db
./..mypy_cache/3.13/cache.13.db
./..mypy_cache/3.13/cache.14.db
./..mypy_cache/3.13/cache.15.db
./..mypy_cache/3.13/cache.2.db
./..mypy_cache/3.13/cache.3.db

```

=== golden label CSVs / gt files ===

```

./output/mahadevpura_singles_golden_features.json
./output/testlarge_short_golden_features.json
./output/wasb_tuned_segments_vs_golden.csv

```

```

=== trajectory CSVs (persisted WASB/tracknet detection) ===
./output/adarsh_avi_full_vs_human.csv
./output/adarsh_avi_vs_human.csv
./output/adarsh_baseline_gated.csv
./output/adarsh_baseline_segments.csv
./output/adarsh_gemini_fullvideo.csv
./output/adarsh_gemini_refined.csv
./output/adarsh_gemini_smoke.csv
./output/adarsh_min5s_segments.csv
./output/adarsh_tuned_segments.csv
./output/gbaaddy_vs_human.csv
./output/kushagra_vs_human.csv
./output/largetestvideo_vs_human.csv
./output/mahadevpura_gemini_refined.csv
./output/mahadevpura_vs_human.csv
./output/testlarge_gemini_fullvideo.csv
./output/testlarge_gemini_rallies.csv
./output/testlarge_vs_human.csv
./output/wasb/GX010125_1080p__89605e80ed01_wasb.csv
./output/wasb/mahadevpura_smoke_180s__05e0e577029d_wasb.csv
./output/wasb_tuned_segments.csv

=== eval_baselines ===
gemini_wrapper_2026-06-10.json
heuristic_lovo_n6.json

```

222. assistant (2026-06-14T03:03:33.921Z)

Promising ? there are **local pre-extracted golden feature files** (`output/*_golden_features.json`), a couple of WASB trajectories, and a manifest. The manifest's trajectory/label paths point outside the repo (`Annotation Setup/...`). Let me check what's actually reachable and understand the golden-feature format + the existing `fusion_golden` driver.

223. user (2026-06-14T03:03:38.228Z)

```

=== external manifest paths exist? ===
adarsh_avi_full_traj.csv
adarsh_avi_traj.csv
adarsh_traj.csv
gbaaddy_traj.csv
kushagra_traj.csv
largetestvideo_traj.csv
mahadevpura_traj.csv
pickleball_traj.csv
Adarsh and Avi.MP4
Adarsh and Avi.rallies.csv
KushagraYashAviFirstVideo.MP4
KushagraYashAviFirstVideo.rallies.csv
KushagraYashAviFirstVideo.report.json
KushagraYashAviFirstVideo.report.md
KushagraYashAviFirstVideo.summary.md
LargeTestVideo.MP4

=== testlarge_short golden_features.json (structure) ===
type <class 'dict'>
keys/len ['name', 'fps', 'frame_width', 'candidates', 'gts']
{
  "name": "testlarge_short",
  "fps": 29.97,
  "frame_width": 1920.0,
  "candidates": [
    {
      "start": 0.06673340006673341,
      "end": 4.8715382048715385,

```

```

"shuttle": {
  "duration": 4.804804804804805,
  "peak_velocity": 0.845096383820243,
  "mean_velocity": 0.19545638143000324,
  "active_ratio": 0.6805555555555556,
  "net_crossings": 5.0
},
"person": {
  "players_in_court": 7.0,
  "two_sided_motion": 0.0,
  "mean_player_motion": 0.0036353791267126423,
  "in_court_ratio": 1.0,
  "shuttle_player_colocation": 0.0
},
"label": 0
},
{
  "start": 8.274941608274942,
  "end": 11.745078411745078,
  "shuttle": {
    "duration": 3.4701368034701368,
    "peak_velocity": 1.321839538670733,
    "mean_velocity": 0.26093356635777015,
    "active_ratio": 0.701923076923076
  }
}

```

224. assistant (2026-06-14T03:05:08.802Z)

This is the "persist once, re-score offline" path working exactly as designed ? ****two golden videos are fully pre-extracted locally**** (`testlarge\_short`, `mahadevpura\_singles`: candidate windows + shuttle/person features + per-candidate labels + golden intervals), and the external trajectories/labels/MP4s exist too. So I can run a ****real**** offline A/B here, no GPU. Let me first probe that the harness imports and the data loads, then build the demo.

225. user (2026-06-14T03:05:18.655Z)

```

loaded testlarge_short | candidates: 14 | golden rallies: 6
CONTROL  agg: {'precision': 0.214, 'recall': 0.5, 'f1': 0.3, 'mean_iou': 0.65}
GateA    agg: {'precision': 0.3, 'recall': 0.5, 'f1': 0.375, 'mean_iou': 0.65}
delta B-A: {'precision': 0.086, 'recall': 0.0, 'f1': 0.075, 'mean_iou': 0.0}

```

226. assistant (2026-06-14T03:05:49.235Z)

Real A/B on a small golden video works: on `testlarge\_short` (14 candidates, 6 rallies), ****Gate-A (net-crossings ?2) lifts precision 0.214?0.30 and F1 0.30?0.375 with recall flat at 0.50**** ? exactly the held-shuttle/low-crossing FP killer behaving as designed. Now let me build the full demo that layers the C1/C4 honest metrics on top, to show what wiring them would actually surface.

227. assistant (2026-06-14T03:06:12.993Z)

"""Offline A/B on a GOLDEN video ? shows the current harness report vs. what the C1/C4 measurement-honesty helpers would add if wired into the default harness.

Runs entirely offline on pre-persisted golden features (no GPU) ? the "persist detection once, re-score forever" path. Honest: prints whatever the real golden data produces.

"""

```
import json
```

```

from backend.eval.experiment import CONTROL, Variant, VideoCandidates, compare, predict
from backend.eval.metrics import score
from backend.eval import segmentation_metrics as sm
from backend.eval import eval_stats as es

```

```
SHUTTLE = ["duration", "peak_velocity", "mean_velocity", "active_ratio", "net_crossings"]
PERSON = ["players_in_court", "two_sided_motion", "mean_player_motion",
          "in_court_ratio", "shuttle_player_colocation"]
```

```
def load(path):
    d = json.load(open(path))
    cand = d["candidates"]
    feats = {}
    for k in SHUTTLE:
        feats[k] = [c["shuttle"][k] for c in cand]
    for k in PERSON:
        feats[k] = [c["person"][k] for c in cand]
    return VideoCandidates(name=d["name"],
                           intervals=[(c["start"], c["end"]) for c in cand],
                           features=feats,
                           gts=[tuple(g) for g in d["gts"]])
```

```
def line(c="-"):
    print(c * 78)
```

```
tl = load("output/testlarge_short_golden_features.json")
mh = load("output/mahadevpura_singles_golden_features.json")
```

```
line("=")
print(f"GOLDEN A/B DEMO ? single small golden video: {tl.name}")
print(f" {len(tl.intervals)} candidate windows, {len(tl.gts)} golden rallies")
line("=")
```

----- Section 1

```
print("\n[1] CURRENT HARNESS A/B ? CONTROL (no gate) vs Gate-A (net_crossings >= 2)")
print(" (a static A/B: no training, so a single golden video is enough)")
gateA = Variant(name="gateA", min_crossings=2)
res = compare([tl], CONTROL, gateA)
a, b = res["variant_a"]["aggregate"], res["variant_b"]["aggregate"]
print(f" {'metric':<12}{'CONTROL':>10}{'Gate-A':>10}{'B - A':>10}")
for m in ("precision", "recall", "f1", "mean_iou"):
    print(f" {m:<12}{a[m]:>10.3f}{b[m]:>10.3f}{res['delta'][m]:>+10.3f}")
preds_ctrl = predict(CONTROL, tl)
preds_gate = predict(gateA, tl)
print(f" segments served: CONTROL={len(preds_ctrl)} Gate-A={len(preds_gate)}
(golden={len(tl.gts)})")
```

----- Section 2

```
print("\n[2] WHAT C4 ADDS ? over-segmentation-aware metrics on the same A/B")
print(" temporal-IoU (above) is nearly blind to fragmentation; F1@k / mAP@tIoU are not")
for name, preds in (("CONTROL", preds_ctrl), ("Gate-A", preds_gate)):
    f1k = sm.f1_at_overlaps(preds, tl.gts)
    ratio = sm.segment_count_ratio(preds, tl.gts)
    print(f" {name:<9} F1@10={f1k[0.1]:.3f} F1@25={f1k[0.25]:.3f} F1@50={f1k[0.5]:.3f}"
          f" seg_count_ratio={ratio:.2f}")
```

proposal-set mAP@tIoU (candidates ranked by peak_velocity as the confidence)

```
scored = [(s, e), pv) for (s, e), pv in zip(tl.intervals, tl.features["peak_velocity"])]
print(f" proposal-set mAP@tIoU[.3-.7] (candidates ranked by peak_velocity) = "
      f"{sm.mean_average_precision(scored, tl.gts):.3f}")
```

----- Section 3

```
print("\n[3] WHAT C1 ADDS ? honest aggregation on a LEARNED A/B (LOVO over the 2 local golden
videos)")
print(" CONTROL vs learned fusion (logistic over shuttle+person features), leave-one-VIDEO-
```

```

out")
learned = Variant(name="learned_fusion", feature_names=SHUTTLE + PERSON)
res2 = compare([tl, mh], CONTROL, learned, honest=True)
pv_a = {p["video"]: p["f1"] for p in res2["variant_a"]["per_video"]}
pv_b = {p["video"]: p["f1"] for p in res2["variant_b"]["per_video"]}
vids = list(pv_a)
print("    per-fold held-out F1:")
for v in vids:
    print(f"        {v:<22} CONTROL={pv_a[v]:.3f}    learned={pv_b[v]:.3f}
B-A={pv_b[v]-pv_a[v]:+.3f}")
a_f1 = [pv_a[v] for v in vids]
b_f1 = [pv_b[v] for v in vids]
print(f"    bare mean F1 (what we report TODAY):  CONTROL={sum(a_f1)/len(a_f1):.3f}  "
      f"learned={sum(b_f1)/len(b_f1):.3f}")
ci_b = es.mean_with_ci(b_f1, seed=0)
delta = es.paired_delta_ci(a_f1, b_f1, seed=0)
st = delta["sign_test"]
print(f"    C1 honest view (what wiring would report):")
print(f"        learned F1 = {ci_b['mean']:.3f}  95% CI [{ci_b['ci_low']:.3f},
{ci_b['ci_high']:.3f}]  (n={int(ci_b['n'])})")
print(f"        mean ?F1 (B-A) = {delta['mean_delta']:+.3f}  CI [{delta['ci_low']:+.3f},
{delta['ci_high']:+.3f}]")
print(f"        excludes 0? {delta['ci_excludes_zero']}")
print(f"        sign test: wins={int(st['wins'])} losses={int(st['losses'])}  "
      f"n_eff={int(st['n_effective'])}  p={st['p_value']:.3f}")
line("=")
print("READ: §1/§2 = real effect on one golden video. §3 shows why a bare LOVO mean is")
print("unsafe at small n ? the CI is wide and the sign test underpowered; that is the")
print("honesty the C1-C4 wiring surfaces BEFORE we trust any 'measured lift'.")
line("=")

```

228. user (2026-06-14T03:06:13.026Z)

File created successfully at: C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/ab_eval_demo.py (file state is current in your context ? no need to Read it back)

229. user (2026-06-14T03:06:17.712Z)

Traceback (most recent call last):

File "C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\ab_eval_demo.py", line 10, in <module>

```

from backend.eval.experiment import CONTROL, Variant, VideoCandidates, compare, predict
ModuleNotFoundError: No module named 'backend'

```

230. assistant (2026-06-14T03:06:25.306Z)

Import path issue (running a script file puts `scratch/` on the path, not the repo root). Re-running with the repo root on `PYTHONPATH`.

231. user (2026-06-14T03:06:29.263Z)

```

=====
GOLDEN A/B DEMO ? single small golden video: testlarge_short
14 candidate windows, 6 golden rallies
=====

```

```

[1] CURRENT HARNESS A/B ? CONTROL (no gate) vs Gate-A (net_crossings >= 2)
(a static A/B: no training, so a single golden video is enough)
metric      CONTROL    Gate-A    B - A
precision    0.214      0.300    +0.086
recall       0.500      0.500    +0.000
f1           0.300      0.375    +0.075
mean_iou     0.650      0.650    +0.000

```

segments served: CONTROL=14 Gate-A=10 (golden=6)

[2] WHAT C4 ADDS ? over-segmentation-aware metrics on the same A/B temporal-IoU (above) is nearly blind to fragmentation; F1@k / mAP@tIoU are not
CONTROL F1@10=0.600 F1@25=0.500 F1@50=0.300 seg_count_ratio=2.33
Gate-A F1@10=0.750 F1@25=0.625 F1@50=0.375 seg_count_ratio=1.67
proposal-set mAP@tIoU[.3-.7] (candidates ranked by peak_velocity) = 0.360

[3] WHAT C1 ADDS ? honest aggregation on a LEARNED A/B (LOVO over the 2 local golden videos)
CONTROL vs learned fusion (logistic over shuttle+person features), leave-one-VIDEO-out
per-fold held-out F1:
testlarge_short CONTROL=0.300 learned=0.300 B-A=+0.000
mahadevpura_singles CONTROL=0.455 learned=0.215 B-A=-0.239
bare mean F1 (what we report TODAY): CONTROL=0.377 learned=0.258
C1 honest view (what wiring would report):
learned F1 = 0.258 95% CI [0.215, 0.300] (n=2)
mean ?F1 (B-A) = -0.120 CI [-0.239, +0.000] excludes 0? False
sign test: wins=0 losses=1 n_eff=1 p=1.000

=====

READ: §1/§2 = real effect on one golden video. §3 shows why a bare LOVO mean is unsafe at small n ? the CI is wide and the sign test underpowered; that is the honesty the C1-C4 wiring surfaces BEFORE we trust any 'measured lift'.

=====

232. assistant (2026-06-14T03:07:30.823Z)

Yes ? ran it **fully offline** on a golden video, no GPU** (the "persist once, re-score" path), on `testlarge\_short` (14 candidate windows, 6 golden rallies). Here's what we actually see, with real numbers:

[1] The A/B itself ? CONTROL vs Gate-A (net-crossings ? 2)

A static A/B, so one golden video is enough:

| metric | CONTROL | Gate-A | ? |
|-----------|---------|-----------|--------|
| precision | 0.214 | **0.300** | +0.086 |
| recall | 0.500 | 0.500 | +0.000 |
| f1 | 0.300 | **0.375** | +0.075 |
| mean_iou | 0.650 | 0.650 | 0.000 |

Gate-A drops 4 of 14 served segments (14?10, golden=6) ? it kills low-crossing false fires, **precision and F1 up, recall untouched.** The held-shuttle/service-feed FP killer working as designed, on real golden data.

[2] What C4 adds ? the over-segmentation that tIoU hid

| variant | F1@10 | F1@25 | F1@50 | **seg-count ratio** |
|---------|-------|-------|-------|---------------------|
| CONTROL | 0.600 | 0.500 | 0.300 | **2.33** |
| Gate-A | 0.750 | 0.625 | 0.375 | **1.67** |

This is the payoff: the **segment-count ratio** shows CONTROL emits 2.33× as many segments as there are rallies** ? heavy fragmentation that the IoU table in [1] barely hinted at ? and Gate-A cuts it to 1.67×. F1@10=0.75 also shows the win is bigger at loose overlap than the strict tIoU-F1 (0.375) alone suggested. Proposal-set mAP@tIoU = 0.360. **This is exactly the effect that's invisible without wiring C4.**

[3] What C1 adds ? why a bare mean would mislead (learned fusion, 2-video LOVO)

| fold | CONTROL F1 | learned F1 | ? |
|-----------------|------------|------------|--------|
| testlarge_short | 0.300 | 0.300 | +0.000 |
| mahadevpura | 0.455 | 0.215 | ?0.239 |

- **Bare mean (what we report today):** CONTROL 0.377 vs learned 0.258 ? "learned is worse, reject it."

- **C1 honest view:** $F1 = 0.120$, 95% CI $[0.239, 0.000]$, sign test $n_{\text{eff}}=1$, $p=1.000$? **inconclusive**. With $n=2$ you cannot conclude anything ? and the wired report says so instead of a misleading point estimate.

Important honesty note on [3]: that learned-fusion result is a demo of the **C1 machinery** at $n=2$, **not** a verdict on learned fusion. The repo's documented $n=6$ result is the opposite (learned 0.583 vs heuristic 0.480) ? learned wins once the corpus is big enough. The $n=2$ underperformance is exactly the small-N artifact C1 exists to flag.

What this tells us about wiring

- **C4 earns its place loudly** ? the segment-count ratio + $F1@k$ surfaced a real fragmentation story plain tIoU buried. Worth wiring.
- **C1 earns its place** ? at our corpus size a single mean is genuinely unsafe; the CI + sign test are the honest gate.
- Both ran against real golden data with zero new infrastructure.

The demo is `scratch/ab_eval_demo.py` (untracked). Natural next steps, your call: **(a)** run the same with C1 stats across **all 6 golden videos** (shuttle-only features are computable here from the external trajectory CSVs ? a real $n=6$ LOVO, the number that actually justifies promotion), or **(b)** I do the default-harness wiring PR so `compare()`/`run_regression.py` emit these by default. Want either?

233. user (2026-06-14T03:45:36.852Z)

Really satisfying to see that the framework is holding up! Well designed! It would be great to do (a), maybe document it as an archived doc for posterity highlighting the first use case of A/B experiment and then following up with 2? No need to wait for my approval, just summarize the (a) outputs and wins/losses and move on to the (b). Makes sense?

234. assistant (2026-06-14T03:46:10.349Z)

Love it ? and you've cleared me to run straight through (a) ? archived write-up ? (b). Let me first look at the existing trajectory?features extraction (`distill_local.py`) so the $n=6$ run uses the **same** windowing as the persisted pipeline rather than a reinvented one.

235. user (2026-06-14T03:46:11.135Z)

```
1      """Distill Gemini (teacher) into a LEARNED local rally classifier (student).
2
3      Threshold-tuning the local knobs hit a low ceiling (see calibrate_local). This goes
4      further: learn "is this WASB candidate a real rally?" from Gemini silver labels,
5      using a few resolution/fps-INVARIANT local features, so the model transfers across
6      videos and runs 100% local at inference (no Gemini, privacy lever intact).
7
8      Flow (all local at runtime):
9      WASB trajectory -> recall-oriented candidate windows (propose many; classifier
10     filters) -> per-candidate features [duration, peak_velocity, mean_velocity,
11     active_ratio, net_crossings] -> LogisticRegression(prob >= threshold) -> merge
12     -> rally windows.
13
14     Training labels come from Gemini full-context labels via the SAME IoU matching the
15     evaluator uses (training.label_candidates). Honest leave-one-video-out: train on the
16     other videos' candidates, predict the held-out one, score vs its Gemini labels, and
17     compare against the threshold baseline. Reuses classifier.py, training.label_candidates,
18     rally_gate, metrics, calibrate_wasb.load_golden_gts, gemini_refine.merge_overlaps.
19
20     Run:
21     python -m backend.eval.distill_local --manifest videos.json --model-out
output/local_rally_model.json
22     """
23     from __future__ import annotations
24
25     import json
26     import argparse
```



```

27     from typing import Dict, List, Optional, Tuple
28
29
30     from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
31     from backend.pipeline.detectors import rally_gate
32     from backend.eval.calibrate_wasb import load_golden_gts
33     from backend.eval.training import label_candidates
34     from backend.eval.classifier import LogisticRegression
35     from backend.eval.metrics import score
36     from backend.eval.gemini_refine import merge_overlaps
37
38     Interval = Tuple[float, float]
39
40     FEATURE_NAMES = ["duration", "peak_velocity", "mean_velocity", "active_ratio",
41 "net_crossings"]
42     # Recall-oriented proposal: deliberately permissive so real rallies are among the
43     # candidates (the classifier removes the junk). (velocity_thresh, merge_gap, min_dur)
44     PROPOSAL = (0.005, 1.0, 0.5)
45     BASELINE_LOCAL = (0.02, 2.0, 2.0) # today's production-ish windowing, no learning
46
47     def build_candidates(trajecory_csv: str, fps: float, frame_width: float,
48                         proposal: Tuple[float, float, float] = PROPOSAL) ->
49 Tuple[List[Interval], List[List[float]]]:
50     """WASB trajectory -> (candidate intervals, fps/resolution-invariant feature
51 rows)."""
52     points = TrackNetRunner.parse_trajectory_csv(trajecory_csv)
53     vt, mg, mwd = proposal
54     wins = TrackNetRunner.trajectory_to_action_windows(
55         points, fps=fps, frame_width=frame_width,
56         velocity_thresh=vt, merge_gap=mg, min_window_duration=mwd)
57     intervals = [(w["start"], w["end"]) for w in wins]
58     gated = rally_gate.gate_windows(points, intervals, fps, min_crossings=0) # for net-
59 crossings
60     X: List[List[float]] = []
61     for w, g in zip(wins, gated):
62         dur = max(1e-6, w["end"] - w["start"])
63         win_frames = max(1.0, dur * fps)
64         X.append([
65             dur, # rally length (sec) ?
66             float(w["peak_velocity"]), # already normalized by
67             float(w["mean_velocity"]),
68             float(w["active_frame_count"]) / win_frames, # active-shuttle ratio ?
69             float(g.crossings), # net crossings (count) ?
70             the rally signal
71         ])
72     return intervals, X
73
74     def baseline_windows(trajecory_csv: str, fps: float, frame_width: float) ->
75 List[Interval]:
76     points = TrackNetRunner.parse_trajectory_csv(trajecory_csv)
77     vt, mg, mwd = BASELINE_LOCAL
78     wins = TrackNetRunner.trajectory_to_action_windows(
79         points, fps=fps, frame_width=frame_width,
80         velocity_thresh=vt, merge_gap=mg, min_window_duration=mwd)
81     return [(w["start"], w["end"]) for w in wins]
82
83     def load_video(spec: Dict, iou: float = 0.5) -> Dict:
84         fps, fw = float(spec["fps"]), float(spec["frame_width"])
85         intervals, X = build_candidates(spec["trajectory"], fps, fw)

```

```

83     gts = load_golden_gts(spec["labels"])
84     y = label_candidates(intervals, gts, iou)
85     return {"name": spec.get("name", spec["trajectory"]), "intervals": intervals, "X":
X,
86             "y": y, "gts": gts, "baseline": baseline_windows(spec["trajectory"], fps,
fw)}}
87
88
89     def predict_rallies(model: LogisticRegression, X: List[List[float]], intervals:
List[Interval],
90                       threshold: float = 0.5) -> List[Interval]:
91         if not intervals:
92             return []
93         proba = model.predict_proba(X)
94         pos = [iv for iv, p in zip(intervals, proba) if p >= threshold]
95         return merge_overlaps(pos, gap_tol=1.0)
96
97
98     def run(specs: List[Dict], iou: float = 0.5, threshold: float = 0.5,
99           model_out: Optional[str] = None) -> dict:
100         videos = [load_video(s, iou) for s in specs]
101         print(f"=== Distilled local rally classifier vs Gemini silver-GT "
102               f"({len(videos)} videos, IoU>={iou}, prob>={threshold}) ===")
103         for v in videos:
104             ceil = score(v["intervals"], v["gts"], iou).recall
105             print(f"[{v['name']}] {len(v['intervals'])} candidates ({sum(v['y'])} pos) | "
106                   f"proposal recall ceiling {ceil:.3f} | {len(v['gts'])} Gemini rallies")
107
108         result: dict = {"videos": [v["name"] for v in videos]}
109         if len(videos) >= 2:
110             print("\n--- HONEST leave-one-video-out: learned vs threshold baseline (held-
out) ---")
111             clf, base = [], []
112             for i, held in enumerate(videos):
113                 others = [v for j, v in enumerate(videos) if j != i]
114                 X = [row for v in others for row in v["X"]]
115                 y = [t for v in others for t in v["y"]]
116                 model = LogisticRegression().fit(X, y, FEATURE_NAMES)
117                 preds = predict_rallies(model, held["X"], held["intervals"], threshold)
118                 sc = score(preds, held["gts"], iou)
119                 bsc = score(held["baseline"], held["gts"], iou)
120                 clf.append(sc)
121                 base.append(bsc)
122                 print(f" held-out {held['name']}: LEARNED P={sc.precision:.3f}
R={sc.recall:.3f} F1={sc.f1:.3f}"
123                       f" | baseline P={bsc.precision:.3f} R={bsc.recall:.3f}
F1={bsc.f1:.3f}")
124             def mean(xs, a):
125                 return sum(getattr(s, a) for s in xs) / len(xs)
126             print(f" MEAN held-out: LEARNED P={mean(clf, 'precision'):.3f}
R={mean(clf, 'recall'):.3f} "
127                   f"F1={mean(clf, 'f1'):.3f} | baseline P={mean(base, 'precision'):.3f} "
128                   f"F1={mean(base, 'f1'):.3f} R={mean(base, 'recall'):.3f}")
129             result["mean_learned_f1"] = mean(clf, "f1")
130             result["mean_baseline_f1"] = mean(base, "f1")
131
132         # Final model on ALL videos (for production / running on a fresh held-out video).
133         Xall = [row for v in videos for row in v["X"]]
134         yall = [t for v in videos for t in v["y"]]
135         final = LogisticRegression().fit(Xall, yall, FEATURE_NAMES)
136         assert final.w is not None # set by fit()
137         weights = dict(zip(FEATURE_NAMES, [round(float(w), 3) for w in final.w]))
138         print(f"\nfinal model standardized weights (sign/size = influence): {weights}")
139         if model_out:
140             final.save(model_out)

```

```

141         print(f"final model (all videos) -> {model_out}")
142         result["model_out"] = model_out
143     return result
144
145
146     def main() -> None:
147         ap = argparse.ArgumentParser(description="Distill Gemini labels into a local rally
classifier.")
148         ap.add_argument("--manifest", required=True,
149                         help="JSON: [{name, trajectory, fps, frame_width, labels}, ...]
(labels = Gemini CSV)")
150         ap.add_argument("--iou", type=float, default=0.5)
151         ap.add_argument("--threshold", type=float, default=0.5, help="classifier probability
cutoff")
152         ap.add_argument("--model-out", default="output/local_rally_model.json")
153         args = ap.parse_args()
154         with open(args.manifest, encoding="utf-8") as f:
155             specs = json.load(f)
156         run(specs, iou=args.iou, threshold=args.threshold, model_out=args.model_out)
157
158
159     if __name__ == "__main__":
160         main()
161

```

236. assistant (2026-06-14T03:47:05.550Z)

`distill\_local.py` already has the exact reusable extraction (`build\_candidates`: trajectory CSV ? windows + `[duration, peak\_velocity, mean\_velocity, active\_ratio, net\_crossings]` via the real pipeline windowing + `rally\_gate`). I'll use that for all 6 so the candidate set is consistent. Let me first probe that it imports and runs offline on one external trajectory.

237. user (2026-06-14T03:47:13.797Z)

```

spec: adarsh_avi_singles
candidates: 172 | features: ['duration', 'peak_velocity', 'mean_velocity', 'active_ratio',
'net_crossings']
golden rallies: 96
sample feature row: [0.851, 1.08, 0.167, 1.02, 0.0]

```

238. assistant (2026-06-14T03:47:49.552Z)

Extraction runs offline on the external golden trajectories. Now I'll write the full n=6 LOVO driver ? one consistent candidate set per video (recall-oriented proposal windowing, same as the persisted pipeline), re-scored three ways (CONTROL / Gate-A / learned-LOVO), with the C1 honest aggregation, C4 segmentation metrics, and an out-of-fold (C3) learned mAP.

239. assistant (2026-06-14T03:48:17.130Z)

```

"""First full A/B experiment on the n=6 golden corpus ? runs entirely offline
(persisted WASB trajectories, no GPU). One consistent candidate set per video
(recall-oriented proposal windowing, the persisted pipeline's own), re-scored three
ways, then graded with the C1 honest aggregation + C4 segmentation metrics + a C3
out-of-fold learned mAP. Prints a report and writes output/ab_experiment_n6.json.

```

```

Honest: prints whatever the real golden data produces. No claims before measurement.
"""

```

```
import json
```

```

from backend.eval.distill_local import FEATURE_NAMES, build_candidates
from backend.eval.calibrate_wasb import load_golden_gts
from backend.eval.training import label_candidates
from backend.eval.classifier import LogisticRegression
from backend.eval.experiment import CONTROL, Variant, VideoCandidates, evaluate_variant

```

```

from backend.eval.metrics import score
from backend.eval import segmentation_metrics as sm
from backend.eval import eval_stats as es

IOU = 0.5
MANIFEST = "output/human_lovo_manifest.json"
OUT = "output/ab_experiment_n6.json"

def load_corpus():
    specs = json.load(open(MANIFEST))
    videos, pooled = [], {"X": [], "Y": [], "group": [], "interval": []}
    for s in specs:
        fps, fw = float(s["fps"]), float(s["frame_width"])
        intervals, X = build_candidates(s["trajectory"], fps, fw)
        gts = load_golden_gts(s["labels"])
        y = label_candidates(intervals, gts, IOU)
        feats = {FEATURE_NAMES[k]: [row[k] for row in X] for k in range(len(FEATURE_NAMES))}
        videos.append(VideoCandidates(name=s["name"], intervals=intervals, features=feats,
gts=gts))
        for iv, row, yi in zip(intervals, X, y):
            pooled["X"].append(row); pooled["Y"].append(yi)
            pooled["group"].append(s["name"]); pooled["interval"].append(iv)
    return videos, pooled

def c4_aggregate(ev, videos):
    """Mean F1@k + mean segment-count ratio across videos for one variant's predictions."""
    by_name = {v.name: v for v in videos}
    f1k = {0.1: [], 0.25: [], 0.5: []}
    ratios = []
    for name, preds in ev["predictions"].items():
        gts = by_name[name].gts
        got = sm.f1_at_overlaps(preds, gts)
        for k in f1k:
            f1k[k].append(got[k])
        ratios.append(sm.segment_count_ratio(preds, gts))
    mean = lambda xs: sum(xs) / len(xs) if xs else 0.0
    return {f"f1@{int(k*100)}": round(mean(v), 3) for k, v in f1k.items()} | \
        {"seg_count_ratio": round(mean(ratios), 2)}

def oof_learned_map(pooled, videos):
    """C3 + C4: out-of-fold learned candidate proba ? mAP@tIoU per video, averaged."""
    def fit_predict(train_idx, test_idx):
        m = LogisticRegression().fit([pooled["X"][i] for i in train_idx],
[pooled["Y"][i] for i in train_idx], FEATURE_NAMES)
        return list(m.predict_proba([pooled["X"][i] for i in test_idx]))
    oof = es.out_of_fold_predictions(pooled["group"], fit_predict)
    by_name = {v.name: v for v in videos}
    maps = []
    for name in by_name:
        scored = [(pooled["interval"][i], float(oof[i]))
            for i in range(len(oof)) if pooled["group"][i] == name and oof[i] is not None]
        maps.append(sm.mean_average_precision(scored, by_name[name].gts))
    return sum(maps) / len(maps) if maps else 0.0

def c1_block(a_f1, b_f1):
    d = es.paired_delta_ci(a_f1, b_f1, seed=0)
    return {
        "mean_a": round(sum(a_f1) / len(a_f1), 3),
        "mean_b": round(sum(b_f1) / len(b_f1), 3),
        "mean_delta": round(d["mean_delta"], 3),
        "ci": [round(d["ci_low"], 3), round(d["ci_high"], 3)],
    }

```

```

        "ci_excludes_zero": d["ci_excludes_zero"],
        "sign": {k: int(v) if k != "p_value" else round(v, 3) for k, v in
d["sign_test"].items()},
    }

def main():
    videos, pooled = load_corpus()
    print("=" * 80)
    print(f"FIRST A/B EXPERIMENT ? n={len(videos)} golden corpus (offline, persisted
trajectories)")
    print("=" * 80)
    for v in videos:
        ceil = score(v.intervals, v.gts, IOU).recall
        print(f" {v.name:<22} {len(v.intervals):>4} candidates {len(v.gts):>3} rallies "
              f"proposal-recall-ceiling={ceil:.3f}")

    control = evaluate_variant(videos, CONTROL, IOU)
    gateA = evaluate_variant(videos, Variant(name="gateA", min_crossings=2), IOU)
    learned = evaluate_variant(videos, Variant(name="learned", feature_names=FEATURE_NAMES),
IOU, honest=True)

    order = [p["video"] for p in control["per_video"]]
    f1 = {ev["variant"]: {p["video"]: p["f1"] for p in ev["per_video"]}}
        for ev in (control, gateA, learned)}

    def col(name):
        return [f1[name][v] for v in order]

    print("\n--- per-video held-out F1 ---")
    print(f" {'video':<22}{'CONTROL':>9}{'GateA':>9}{'learned':>9}")
    for v in order:
        print(f"
{v:<22}{f1['CONTROL'][v]:>9.3f}{f1['gateA'][v]:>9.3f}{f1['learned'][v]:>9.3f}")

    ab_gate = c1_block(col("CONTROL"), col("gateA"))
    ab_learned = c1_block(col("CONTROL"), col("learned"))
    print("\n--- A/B 1: CONTROL vs Gate-A (static) ? C1 honest aggregation ---")
    print(f" meanF1 {ab_gate['mean_a']} -> {ab_gate['mean_b']} ?F1={ab_gate['mean_delta']:+} "
          f"CI{ab_gate['ci']} excl0={ab_gate['ci_excludes_zero']} sign={ab_gate['sign']}")
    print("--- A/B 2: CONTROL vs learned fusion (LOVO) ? C1 honest aggregation ---")
    print(f" meanF1 {ab_learned['mean_a']} -> {ab_learned['mean_b']}
?F1={ab_learned['mean_delta']:+} "
          f"CI{ab_learned['ci']} excl0={ab_learned['ci_excludes_zero']}
sign={ab_learned['sign']}")

    c4 = {name: c4_aggregate(ev, videos)
          for name, ev in (("CONTROL", control), ("gateA", gateA), ("learned", learned))}
    print("\n--- C4 over-segmentation-aware metrics (mean across videos) ---")
    for name in ("CONTROL", "gateA", "learned"):
        print(f" {name:<9} {c4[name]}")
    oof_map = oof_learned_map(pooled, videos)
    print(f" learned out-of-fold (C3) mAP@tIoU[.3-.7] = {oof_map:.3f}")

    result = {
        "corpus": order, "n": len(videos), "iou": IOU,
        "aggregate_f1": {k: round(sum(col(k)) / len(order), 3) for k in ("CONTROL", "gateA",
"learned")},
        "per_video_f1": f1,
        "ab_control_vs_gateA": ab_gate,
        "ab_control_vs_learned": ab_learned,
        "c4_metrics": c4,
        "learned_oof_map": round(oof_map, 3),
    }
    json.dump(result, open(OUT, "w"), indent=2)

```

```

print(f"\nwrote {OUT}")

if __name__ == "__main__":
    main()

```

240. user (2026-06-14T03:48:17.151Z)

File created successfully at: C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/ab_n6_lovo.py (file state is current in your context ? no need to Read it back)

241. user (2026-06-14T03:48:27.996Z)

```

=====
FIRST A/B EXPERIMENT ? n=6 golden corpus (offline, persisted trajectories)
=====
adarsh_avi_singles      172 candidates   96 rallies   proposal-recall-ceiling=0.833
kushagra_singles        20 candidates   32 rallies   proposal-recall-ceiling=0.125
largetest_doubles       86 candidates   49 rallies   proposal-recall-ceiling=0.816
gbaaddy                 82 candidates   48 rallies   proposal-recall-ceiling=0.312
testlarge_short         14 candidates    6 rallies   proposal-recall-ceiling=0.500
mahadevpura_singles     57 candidates   31 rallies   proposal-recall-ceiling=0.645

--- per-video held-out F1 ---
video          CONTROL   GateA   learned
Traceback (most recent call last):
  File "C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\ab_n6_lovo.py", line 142, in
<module>
    main()
    ~~~~^^
  File "C:\Users\avidu\Projects\badminton-highlight-indexer\scratch\ab_n6_lovo.py", line 109, in
main
    print(f"   {v:<22}{f1['CONTROL'][v]:>9.3f}{f1['gateA'][v]:>9.3f}{f1['learned'][v]:>9.3f}")
          ~~~~~^
KeyError: 'CONTROL'

```

242. assistant (2026-06-14T03:48:37.169Z)

Bug ? `CONTROL.name` is `"control"`, not my display key `"CONTROL"`. Let me key the F1 dict explicitly.

243. user (2026-06-14T03:48:39.377Z)

The file C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/ab_n6_lovo.py has been updated successfully. (file state is current in your context ? no need to Read it back)

244. user (2026-06-14T03:48:53.897Z)

```

=====
FIRST A/B EXPERIMENT ? n=6 golden corpus (offline, persisted trajectories)
=====
adarsh_avi_singles      172 candidates   96 rallies   proposal-recall-ceiling=0.833
kushagra_singles        20 candidates   32 rallies   proposal-recall-ceiling=0.125
largetest_doubles       86 candidates   49 rallies   proposal-recall-ceiling=0.816
gbaaddy                 82 candidates   48 rallies   proposal-recall-ceiling=0.312
testlarge_short         14 candidates    6 rallies   proposal-recall-ceiling=0.500
mahadevpura_singles     57 candidates   31 rallies   proposal-recall-ceiling=0.645

--- per-video held-out F1 ---
video          CONTROL   GateA   learned
adarsh_avi_singles      0.597    0.722    0.613
kushagra_singles        0.154    0.163    0.163
largetest_doubles       0.593    0.667    0.535
gbaaddy                0.231    0.274    0.222
testlarge_short         0.300    0.375    0.000

```

```

mahadevpura_singles      0.455      0.571      0.308

--- A/B 1: CONTROL vs Gate-A (static) ? C1 honest aggregation ---
  meanF1 0.388 -> 0.462 ?F1=+0.074 CI[0.042, 0.104] excl0=True  sign={'wins': 6, 'losses': 0,
'ties': 0, 'n_effective': 6, 'p_value': 0.031}
--- A/B 2: CONTROL vs learned fusion (LOVO) ? C1 honest aggregation ---
  meanF1 0.388 -> 0.307 ?F1=-0.081 CI[-0.173, -0.002] excl0=True  sign={'wins': 2, 'losses': 4,
'ties': 0, 'n_effective': 6, 'p_value': 0.688}

--- C4 over-segmentation-aware metrics (mean across videos) ---
  CONTROL   {'f1@10': 0.645, 'f1@25': 0.603, 'f1@50': 0.388, 'seg_count_ratio': 1.68}
  gateA     {'f1@10': 0.748, 'f1@25': 0.698, 'f1@50': 0.462, 'seg_count_ratio': 1.01}
  learned   {'f1@10': 0.564, 'f1@25': 0.509, 'f1@50': 0.307, 'seg_count_ratio': 0.89}
  learned out-of-fold (C3) mAP@tIoU[.3-.7] = 0.245

wrote output/ab_experiment_n6.json

```

245. assistant (2026-06-14T03:50:12.231Z)

The first n=6 A/B experiment produced a genuinely informative, honest result. Let me capture the headline before packaging it:

```

- **Corpus + proposal recall ceilings** expose the known recall trap (kushagra 0.125, gbaaddy
0.312 ? multi-court).
- **A/B-1 CONTROL vs Gate-A:** F1 0.388?0.462, **wins all 6 videos, sign-test p=0.031, ?F1 CI
[+0.042, +0.104] excludes 0** ? a real, statistically-supported win; seg-count-ratio
1.68?1.01** (over-segmentation nearly eliminated).
- **A/B-2 CONTROL vs learned fusion (LOVO):** F1 0.388?0.307, 2W/4L, **does not clear the bar**
(testlarge zeroed out ? a threshold/calibration failure, exactly what C2 targets); OOF mAP
0.245.

```

The harness correctly **promoted a real win and refused a non-win** ? the C1 stats prevented a wrong call in both directions. Let me check the archive index format, then write the posterity doc.

246. user (2026-06-14T03:50:13.030Z)

```

1      # quality-iterations/
2
3      This directory holds completed phases of quality improvement work.
4
5      As we receive new labeled videos and iterate on model accuracy, recall/precision, audio
fusion, candidate quality, etc., each major iteration or batch of improvements will be captured
here once it reaches a terminal state (DONE / COMPLETED / REJECTED).
6
7      See `../README.md` (under archives/) for the overall archival policy.
8
9      ## Organization
10     - Dated subdirectories, e.g. `2026-06-initial/`
11     - Each iteration should include:
12       - A README.md with STATUS, trigger (new data, specific goals), key changes, and links
to relevant checkpoints under `../checkpoints/`.
13       - References to updated eval baselines, regression results, report enhancements, etc.
14
15     This keeps the live `docs/` clean while preserving a rich historical trail for
regression detection across quality work.
16
17     ## Iterations
18     - **2026-06-initial/`** ? SEED. Opened the quality phase ahead of the awaited labeled-
video batch (status: in progress; superseded for first results by the iteration below).
19     - **`2026-06-human-gt-tuning/`** ? **COMPLETED (2026-06-08).** First eval against
**human** ground truth on a partially-labeled video: built the partial-label eval harness +
labeled-window methodology, produced the first trustworthy F1 (baseline 0.650 ? tuned 0.742,
*fit-on-self*), per-rally failure analysis, a config sweep, and ? via adversarial review ? found
+ fixed a real ingestion bug (negative-start import abort). Cross-repo: `sports-data-collector`

```

PR #13.

20 - `**`2026-06-multivideo-lovo/`** ? **COMPLETED (2026-06-08 ? extended to n=6, 2026-06-09).**` Built a `**6-match**` human corpus (3 singles incl. Mahadevpura, 1 doubles, 2 multi-court; 3 cameras/frame-rates) + cross-video `**leave-one-video-out**`. Key finding: the unsupervised windowing heuristic is `**footage-fragile**` ? it fails on `**multi-court**` footage (background games; Kushagra 0.16 / gBaaddy 0.28) where the shuttle signal is present (97% of rallies have motion) but rally-vs-dead-time isn't velocity-separable. `**Prototyped a learned per-frame segmenter**` (``backend/eval/rally_seq_proto.py``): held-out AUC `**0.83?0.92` on every match, decoupled from the multi-court confound`**`. `**The learned model OVERTOOK the heuristic at n=3 and holds ~+0.10**`: n=6 LOVO mean `**learned 0.583 vs heuristic 0.480**`. The `**"contrast" metric**` (in-rally÷dead-time shuttle visibility) predicts heuristic F1 monotonically across 5 matches. Data-backed green light for ``../../research/OWNED_MODEL_TRAINING_STUDY.md``.
21

247. assistant (2026-06-14T03:50:48.697Z)

Format clear. Important honesty note I'll bake into the doc: my A/B-2 "learned" is the `*window-level* 5-feature logistic (distill_local-style)`, `**not**` the per-frame prototype (``rally_seq_proto.py``) that scored the documented 0.583 ? so this isn't a contradiction, and `testlarge?0.000` is a threshold/calibration artifact (the C2 motivation). Let me set up a worktree off master and write the posterity doc.

248. user (2026-06-14T03:50:53.549Z)

d7b1f51 feat(eval): honest ?F1 ? bootstrap CI + paired sign test on the fusion lift (#127)
HEAD is now at d7b1f51 feat(eval): honest ?F1 ? bootstrap CI + paired sign test on the fusion lift (#127)
dir ready

249. assistant (2026-06-14T03:51:03.282Z)

Note: master HEAD is now `**#127 "honest ?F1 ? bootstrap CI + paired sign test on the fusion lift"*` ? a sibling has started wiring honest stats into the fusion comparison, which overlaps with (b). I'll check that before doing (b) to avoid duplication. First, the posterity doc:

250. assistant (2026-06-14T03:51:40.790Z)

First end-to-end A/B experiment on the golden corpus (n=6)

`**STATUS: COMPLETED (2026-06-14).**` The first time the experiment harness (`[EXPERIMENT_HARNESS.md](../../EXPERIMENT_HARNESS.md)`) + the measurement-honesty helpers (course corrections C1?C4, `[ANALYZERS_AND_RECONCILERS.md](../../ANALYZERS_AND_RECONCILERS.md)`) §13) were run end-to-end on real golden data ? a posterity record of the framework's first real use.

`**Trigger.**` After landing the signal-fusion framework doc + §13 literature review and the additive C1?C4 helpers (``backend/eval/eval_stats.py``, ``score_calibration.py``, ``segmentation_metrics.py``, PR #124), the question was: `*does the framework actually hold up on a golden video?* This is the validation run.`

Method (100% offline ? the "persist detection once, re-score forever" path)

- `**No GPU.**` Re-scored the `**persisted WASB shuttle trajectories**` of the 6 golden matches (``output/human_lovo_manifest.json``); detection was never re-run.
- `**One consistent candidate set per video**` via ``distill_local.build_candidates`` (recall-oriented proposal windowing ``velocity_thresh=0.005, merge_gap=1.0, min_dur=0.5``) ? per-window shuttle features ``[duration, peak_velocity, mean_velocity, active_ratio, net_crossings]``; golden labels via ``calibrate_wasb.load_golden_gts``; per-candidate y via the evaluator's own IoU matcher.
- `**Three variants re-scored on that same set**` (``backend/eval/experiment.py``):
 - `**CONTROL**` ? merge all candidates (no filter).

- **Gate-A** ? keep candidates with `net_crossings ? 2`, then merge (the held-shuttle / service-feed structural FP killer; `rally_gate`).
- **learned** ? a window-level logistic regression over the 5 shuttle features, **leave-one-VIDEO-out** (train on the other 5 matches, predict held-out), prob ? 0.5, merge.
- **Graded with:** temporal-IoU F1 (per video), **C1** honest aggregation (`mean_with_ci`, `paired_delta_ci`, exact `paired_sign_test`), **C4** segmentation metrics (`f1_at_overlaps`, `segment_count_ratio`), and a **C3** out-of-fold learned mAP@tIoU.
- IoU 0.5. Repro: `scratch/ab_n6_lovo.py` (machine-side; reads the gitignored manifest + the external golden trajectories/labels); raw result `output/ab_experiment_n6.json`.

Corpus

| video | candidates | golden rallies | proposal recall | ceiling |
|--------------------------------|------------|----------------|-----------------|---------|
| adarsh_avi_singles | 172 | 96 | 0.833 | |
| kushagra_singles (multi-court) | 20 | 32 | 0.125 | |
| largetest_doubles | 86 | 49 | 0.816 | |
| gbaaddy (multi-court) | 82 | 48 | 0.312 | |
| testlarge_short | 14 | 6 | 0.500 | |
| mahadevpura_singles | 57 | 31 | 0.645 | |

The multi-court matches (kushagra, gbaaddy) have **low proposal-recall ceilings** ? the decision stage can never recover rallies the windowing missed (the proposal-recall ceiling pitfall, §13). Their absolute F1 is windowing-capped, not a decision failure.

Results

Per-video held-out F1 (IoU 0.5):

| video | CONTROL | Gate-A | learned (LOVO) |
|---------------------|--------------|--------------|----------------|
| adarsh_avi_singles | 0.597 | 0.722 | 0.613 |
| kushagra_singles | 0.154 | 0.163 | 0.163 |
| largetest_doubles | 0.593 | 0.667 | 0.535 |
| gbaaddy | 0.231 | 0.274 | 0.222 |
| testlarge_short | 0.300 | 0.375 | 0.000 |
| mahadevpura_singles | 0.455 | 0.571 | 0.308 |
| mean | 0.388 | 0.462 | 0.307 |

A/B 1 ? CONTROL vs Gate-A (static). ?F1 **+0.074**, 95% CI **[+0.042, +0.104]** (excludes 0), sign test **6 wins / 0 losses**, $p = 0.031$. ? **A real, statistically-supported win** ? Gate-A improves F1 on **every** golden match.

A/B 2 ? CONTROL vs window-level learned logistic (LOVO). ?F1 **0.081**, CI **[0.173, 0.002]**, sign test **2 wins / 4 losses**, $p = 0.688$. ? **Does not clear the bar** ? this variant is **worse** than CONTROL on this setup (and the mean is dragged by two big per-video losses; the sign test is inconclusive ? the two honest views disagreeing is itself a reason to report both).

C4 ? over-segmentation-aware metrics (mean across videos):

| variant | F1@10 | F1@25 | F1@50 | segment-count ratio |
|---------------|--------------|--------------|--------------|---------------------|
| CONTROL | 0.645 | 0.603 | 0.388 | 1.68 |
| Gate-A | 0.748 | 0.698 | 0.462 | 1.01 |
| learned | 0.564 | 0.509 | 0.307 | 0.89 |

Gate-A drives the **segment-count ratio 1.68 ? 1.01** ? it nearly eliminates the over-

segmentation that
plain temporal-IoU barely reflected. The learned variant's 0.89 is **under*-segmentation* (it drops real rallies). ***C3 out-of-fold learned mAP@tIoU[.3?.7] = 0.245*** (weak ranking ? corroborates the loss).

Wins / losses / lessons

- ? ***Gate-A (net-crossings ? 2) is a genuine win*** ? +0.074 mean F1, 6/6 matches, p = 0.031, CI clears 0,
over-segmentation 1.68?1.01. Promote-worthy by the honest gate. (Confirms the held-shuttle owner-bug
killer on the whole corpus, offline, no GPU.)
- ? ***The window-level 5-feature logistic (threshold 0.5) is not a win here*** ? ?0.081, mixed (2W/4L).
`testlarge\_short ? 0.000` (the model predicted **zero** rallies) is a ***threshold/calibration artifact***, a
textbook motivation for ***C2 (calibrate cue confidence)*** and tuning the threshold **inside** the LOVO loop.
- ?? ***Do not conflate this with the documented learned winner.*** The
[`2026-06-multivideo-lovo`](../2026-06-multivideo-lovo/README.md) result (learned 0.583 vs heuristic
0.480) is a ***per-frame prototype*** (`rally\_seq\_proto.py`), a **different, stronger** model than this
window-level logistic. This experiment compared a **weaker** learned path against a **strong** CONTROL
(merge-all of a recall-oriented proposal) ? a different matchup. The takeaway is methodological, not
"learned fusion fails."
- ? ***The methodology itself worked.*** C1 gave a confident ***YES*** on Gate-A **and** a clear ***NO*** on the
logistic ? preventing a wrong promotion in ***both*** directions. A bare-mean report would have hidden
Gate-A's significance and might have over-read the logistic's mean. C4 surfaced the over-segmentation
story tIoU buried. This is exactly the framework holding up.
- ? ***Multi-court is windowing-capped*** ? kushagra/gbaaddy absolute F1 is bounded by the 0.125/0.312
proposal-recall ceiling; the decision layer can't fix what the proposal missed (measure AR@AN separately).

Significance

First real A/B on the golden corpus ? the harness + C1?C4 produced an ***honest, promotable finding*
(Gate-A)**** and an ***honest rejection (this logistic)***, entirely offline. This is the empirical backbone the
paper aim ([`paper-coauthor-aim`]) needs ? honest measured results only. Gate-A becomes a
candidate for the
no-regression promotion gate; the logistic's failure mode is the concrete task list for C2 +
threshold-in-loop.

Links: [EXPERIMENT_HARNESS.md](../.../EXPERIMENT_HARNESS.md) ·
[ANALYZERS_AND_RECONCILERS.md](../.../ANALYZERS_AND_RECONCILERS.md) §13 (C1?C4) ·
[QUALITY_ITERATIONS.md](../.../QUALITY_ITERATIONS.md) · predecessor
[`2026-06-multivideo-lovo`](../2026-06-multivideo-lovo/README.md).

251. user (2026-06-14T03:51:40.821Z)

File created successfully at: C:/Users/avidu/Projects/bhi-ab-wt/docs/archives/quality-iterations/2026-06-first-ab-experiment/README.md (file state is current in your context ? no need to Read it back)

252. assistant (2026-06-14T03:51:47.205Z)

Now the archive index entry and a live-log pointer.

253. user (2026-06-14T03:52:01.370Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

254. user (2026-06-14T03:52:06.659Z)

```
17     ## Iterations
18     - **`2026-06-initial/`** ? SEED. Opened the quality phase ahead of the awaited labeled-
video batch (status: in progress; superseded for first results by the iteration below).
19     - **`2026-06-human-gt-tuning/`** ? **COMPLETED (2026-06-08).** First eval against
**human** ground truth on a partially-labeled video: built the partial-label eval harness +
labeled-window methodology, produced the first trustworthy F1 (baseline 0.650 ? tuned 0.742,
*fit-on-self*), per-rally failure analysis, a config sweep, and ? via adversarial review ? found
+ fixed a real ingestion bug (negative-start import abort). Cross-repo: `sports-data-collector`
PR #13.
20     - **`2026-06-multivideo-lovo/`** ? **COMPLETED (2026-06-08 ? extended to n=6,
2026-06-09).** Built a **6-match** human corpus (3 singles incl. Mahadevpura, 1 doubles, 2
multi-court; 3 cameras/frame-rates) + cross-video **leave-one-video-out**. Key finding: the
unsupervised windowing heuristic is **footage-fragile** ? it fails on **multi-court** footage
(background games; Kushagra 0.16 / gBaaddy 0.28) where the shuttle signal is present (97% of
rallies have motion) but rally-vs-dead-time isn't velocity-separable. **Prototyped a learned
per-frame segmenter** (`backend/eval/rally_seq_proto.py`): held-out AUC **0.83?0.92 on every
match, decoupled from the multi-court confound**. **The learned model OVERTOOK the heuristic at
n=3 and holds ~+0.10**: n=6 LOVO mean **learned 0.583 vs heuristic 0.480**. The **"contrast"
metric** (in-rally÷dead-time shuttle visibility) predicts heuristic F1 monotonically across 5
matches. Data-backed green light for `../../research/OWNED_MODEL_TRAINING_STUDY.md`.
21
```

255. user (2026-06-14T03:52:15.515Z)

The file C:/Users/avidu/Projects/bhi-ab-wt/docs/archives/quality-iterations/README.md has been updated successfully. (file state is current in your context ? no need to Read it back)

256. assistant (2026-06-14T03:52:31.983Z)

Now a concise live-log pointer in `QUALITY\_ITERATIONS.md`. Let me read its current live-track header in the worktree (siblings have edited it).

257. user (2026-06-14T03:52:33.145Z)

```
40
41     ## Current track (live) ? newest first
42
43     ### Eval-honesty helpers (course corrections C1?C4) ? additive, default-OFF
*(2026-06-14)*
44     **Why.** A multi-agent literature review (folded into
45     [ANALYZERS_AND_RECONCILERS.md](ANALYZERS_AND_RECONCILERS.md) §13) confirmed the
framework's *skeleton* is
46     standard prior art (Temporal Action Localization/Segmentation + late/score-level fusion
? cite, don't claim),
47     and flagged that our **measurement** is walking into known small-N traps that would make
"promote only on
```

258. user (2026-06-14T03:52:43.806Z)

The file C:/Users/avidu/Projects/bhi-ab-wt/docs/QUALITY_ITERATIONS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

259. assistant (2026-06-14T03:52:49.534Z)

Now commit the posterity doc (3 files), push, and self-merge on green CI (docs-only).

260. user (2026-06-14T03:53:06.773Z)

warning: in the working copy of 'docs/archives/quality-iterations/2026-06-first-ab-experiment/README.md', LF will be replaced by CRLF the next time Git touches it

=== staged ===

M docs/QUALITY_ITERATIONS.md

A docs/archives/quality-iterations/2026-06-first-ab-experiment/README.md

M docs/archives/quality-iterations/README.md

To <https://github.com/avidullu/badminton-highlight-indexer.git>

* [new branch] docs/first-ab-experiment -> docs/first-ab-experiment

tail:

Line |

38 | Generated with Claude Code." 2>&1 | tail -2

|

~~~~~

| The term 'tail' is

not recognized as a name of a cmdlet, function, script file, or executable program.

Check the spelling of the

name, or if a path was included, verify that the path is correct and try again.

#### 261. assistant (2026-06-14T03:53:15.165Z)

The push worked but my `| tail` (not a PowerShell cmdlet) swallowed the `gh pr create` output.

Let me check whether the PR was created.

#### 262. user (2026-06-14T03:53:16.832Z)

(Bash completed with no output)

#### 263. assistant (2026-06-14T03:53:22.665Z)

PR wasn't created. Let me create it via Bash (proper piping).

#### 264. user (2026-06-14T03:53:33.105Z)

<https://github.com/avidullu/badminton-highlight-indexer/pull/128>

#### 265. assistant (2026-06-14T03:53:38.780Z)

PR [#128](<https://github.com/avidullu/badminton-highlight-indexer/pull/128>). Watching CI, then self-merging (docs-only).

#### 266. user (2026-06-14T03:55:23.359Z)

Ruff lint pass 10s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487606576/job/81246713201>

Syntax sweep + import smoke (no GPU stack) pass 16s

<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487606576/job/81246713210>

Full test suite (offline, mocked) pass 1m31s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487606576/job/81246713203>

Mypy (lenient baseline) pass 29s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487606576/job/81246713199>

Ruff lint pass 10s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487606576/job/81246713201>

Syntax sweep + import smoke (no GPU stack) pass 16s

<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487606576/job/81246713210>  
exit:0

#### 267. assistant (2026-06-14T03:55:26.045Z)

All four lanes green. Merging #128 and cleaning up the worktree.

#### 268. user (2026-06-14T03:55:32.460Z)

{"mergedAt":"2026-06-14T03:55:31Z","state":"MERGED"}

worktree removed  
Deleted branch docs/first-ab-experiment (was d657226).

#### 269. assistant (2026-06-14T03:55:47.112Z)

Archive doc merged (#128). Before doing (b), I must check what sibling **#127** "honest ?F1 ? bootstrap CI + paired sign test" already wired ? it overlaps directly with C1 and (b). Let me investigate to scope (b) without duplicating.

#### 270. user (2026-06-14T03:55:53.388Z)

```
=== #127 changed files ===
d7b1f51 feat(eval): honest ?F1 ? bootstrap CI + paired sign test on the fusion lift (#127)
  backend/eval/fusion_audio.py    | 6 +++++
  backend/eval/fusion_compare.py | 16 ++++++
  tests/test_fusion_compare.py    | 15 ++++++
  3 files changed, 37 insertions(+)

=== who consumes my eval_stats / segmentation_metrics now? ===
backend/eval/eval_stats.py
backend/eval/fusion_compare.py

=== is compare() in experiment.py already emitting CIs? ===
```

#### 271. assistant (2026-06-14T03:56:46.956Z)

Good ? that de-risks (b): sibling **#127** already wired C1 (bootstrap CI + sign test) into ``fusion_compare.py``, and it reuses my ``eval_stats`` (no duplication). But the **core** ``experiment.compare()`` still reports bare means, and **C4** (``segmentation_metrics``) isn't wired anywhere. So (b) is now precisely scoped: wire C1 + C4 into ``experiment.compare()`` as **additive** fields (existing numbers unchanged, honest numbers added). Let me set up a worktree and check the tests won't break.

#### 272. user (2026-06-14T03:56:52.182Z)

```
HEAD is now at 8398e52 docs(archive): first end-to-end A/B experiment on the golden corpus (n=6)
(#128)
=== compare() result-shape assertions in the test? ===
121:  pv = res["per_video"][0]
125:  assert res["aggregate"]["recall"] == pytest.approx(1.0)
126:  assert not res["honest_lovo"]
132:  assert res["honest_lovo"] is True
133:  assert res["aggregate"]["f1"] == pytest.approx(1.0)      # held-out F1 perfect on
separable data
151:  res = compare(videos, CONTROL, Variant(name="gateA", min_crossings=2), iou=0.5)
152:  assert res["variant_b"]["aggregate"]["precision"] == pytest.approx(1.0)
153:  assert res["delta"]["precision"] > 0                      # B (gate) beats A (control)
on precision
154:  assert res["delta"]["f1"] > 0
155:  assert res["label_set_hash"]                               # fingerprint present
156:  assert res["num_videos"] == 1
157:  assert len(res["per_video_delta"]) == 1
162:  compare([], CONTROL, CONTROL)
170:  assert res["num_a"] == 3 and res["num_b"] == 2
171:  assert res["num_agreed"] == 2
172:  assert res["a_only"] == [(12.0, 13.0)]                     # A fired, B suppressed
173:  assert res["b_only"] == []                                  # B is a subset of A here
174:  assert res["mean_agree_iou"] == pytest.approx(1.0)
193:  assert res["num_a"] == 2                                    # control keeps both
194:  assert res["num_b"] == 1                                    # model drops the x=0 window
195:  assert res["a_only"] == [(40.0, 41.0)]
202:  res = compare([_held_shuttle_video()], CONTROL, Variant(name="gateA", min_crossings=2))
225:  assert vc.features["net_crossings"] == [4.0, 0.0]
226:  assert vc.features["duration"] == [10.0, 1.0]              # derived from bounds
```

### 273. assistant (2026-06-14T03:57:03.211Z)

Tests check specific keys only (no exact-dict equality) ? so additive fields are safe. Let me read the current `experiment.py` regions I'll touch (it may have shifted since I last read it).

### 274. user (2026-06-14T03:57:03.807Z)

```
44     from datetime import datetime, timezone
45     from hashlib import sha1
46     from typing import Any, Dict, List, Optional, Sequence
47
48     from backend.eval.calibration import pct_improvement
49     from backend.eval.classifier import LogisticRegression
50     from backend.eval.gemini_refine import merge_overlaps
51     from backend.eval.metrics import Interval, match_intervals, score
52     from backend.eval.regression import gt_hash
53     from backend.eval.training import label_candidates
54
55     logger = logging.getLogger(__name__)
56
57     # Where a comparison run appends one reproducible record. Tracked location (NOT under
58     # the gitignored output/) so the leaderboard survives a clean checkout, mirroring
59     # regression.DEFAULT_BASELINE_PATH.
60     DEFAULT_LEADERBOARD_PATH = os.path.join("eval_baselines", "experiment_leaderboard.json")
61     _METRICS = ("precision", "recall", "f1", "mean_iou")
```

### 275. user (2026-06-14T03:57:08.313Z)

```
303     def compare(videos: Sequence[VideoCandidates], variant_a: Variant, variant_b: Variant,
304                 iou: float = 0.5, honest: bool = True) -> Dict[str, Any]:
305         """Offline labeled A/B (`EXPERIMENT_HARNESS.md` §5.1). Scores both variants on the
same
306         golden corpus and reports the **B ? A** deltas, per video and in aggregate.
307
308         The deltas use the existing `metrics.score` (per video) +
`calibration.pct_improvement`;
309         learned variants are scored LOVO-honestly. The result is a self-contained record fit
for
310         `record_experiment` ? variant specs + hashes + the label-set fingerprint travel with
it.
311         """
312         if len(videos) < 1:
313             raise ExperimentError("compare needs at least one labeled video")
314         eval_a = evaluate_variant(videos, variant_a, iou, honest)
315         eval_b = evaluate_variant(videos, variant_b, iou, honest)
316
317         delta = {m: eval_b["aggregate"][m] - eval_a["aggregate"][m] for m in _METRICS}
318         delta_pct = {m: pct_improvement(eval_a["aggregate"][m], eval_b["aggregate"][m]) for
m in _METRICS}
319
320         by_video_a = {p["video"]: p for p in eval_a["per_video"]}
321         per_video_delta = [
322             {"video": p["video"], **{m: p[m] - by_video_a[p["video"]][m] for m in _METRICS}}
323             for p in eval_b["per_video"]]
324         ]
325
326         # One fingerprint over the whole label set ? detects "the deltas were measured
against
327         # different labels" the same way regression.gt_hash guards the no-regression
baseline.
328         all_gts: List[Interval] = [iv for vc in videos for iv in (vc.gts or [])]
329
330         return {
331             "iou": iou,
332             "label_set_hash": gt_hash(all_gts),
```

```

333         "num_videos": len(videos),
334         "variant_a": eval_a,
335         "variant_b": eval_b,
336         "delta": delta,
337         "delta_pct": delta_pct,
338         "per_video_delta": per_video_delta,
339     }
340
341
342     # ----- SxS on unlabeled video

```

## 276. assistant (2026-06-14T03:57:26.483Z)

Now let me read `record\_experiment` so I can additively record the honest delta too.

## 277. user (2026-06-14T03:57:27.016Z)

```

395     def record_experiment(result: Dict[str, Any], path: str = DEFAULT_LEADERBOARD_PATH,
396                          timestamp: Optional[str] = None) -> Dict[str, Any]:
397         """Append a compact, reproducible record of a `compare()` result to the JSON
398         leaderboard (`EXPERIMENT_HARNESS.md` §6: experiments are records). Generalises
399         `regression_baseline.json`; deliberately the size of a baseline file, not a service
400         (§9). Returns the row written. ``timestamp`` is injectable for deterministic tests.
401         """
402         row: Dict[str, Any] = {
403             "timestamp": timestamp or datetime.now(timezone.utc).isoformat(),
404             "iou": result.get("iou"),
405             "label_set_hash": result.get("label_set_hash"),
406             "num_videos": result.get("num_videos"),
407             "variant_a": {"name": result["variant_a"]["variant"],
408                          "spec_hash": result["variant_a"]["spec_hash"],
409                          "aggregate": result["variant_a"]["aggregate"]},
410             "variant_b": {"name": result["variant_b"]["variant"],
411                          "spec_hash": result["variant_b"]["spec_hash"],
412                          "aggregate": result["variant_b"]["aggregate"]},
413             "delta": result.get("delta"),
414         }
415         entries: List[Dict[str, Any]] = []
416         if os.path.isfile(path):
417             with open(path, encoding="utf-8") as f:
418                 entries = json.load(f)
419         entries.append(row)
420         os.makedirs(os.path.dirname(path) or ".", exist_ok=True)
421         with open(path, "w", encoding="utf-8") as f:
422             json.dump(entries, f, indent=2)
423         logger.info("recorded experiment %s vs %s (?f1=%3f) -> %s",
424                   row["variant_a"]["name"], row["variant_b"]["name"],
425                   (result.get("delta") or {}).get("f1", 0.0), path)
426         return row
427
428
429     def load_video_candidates(db: Any, video_id: str, *, detector: Optional[str] = None,

```

## 278. assistant (2026-06-14T03:57:32.480Z)

Now the additive edits to `experiment.py` ? imports, helpers, the `compare()` fields, and the leaderboard row.

## 279. user (2026-06-14T03:57:35.084Z)

The file C:/Users/avidu/Projects/bhi-wire-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

## 280. user (2026-06-14T03:57:56.371Z)

The file C:/Users/avidu/Projects/bhi-wire-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**281. user (2026-06-14T03:58:03.855Z)**

The file C:/Users/avidu/Projects/bhi-wire-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**282. assistant (2026-06-14T03:58:09.963Z)**

Now additively record the honest ?F1 in the leaderboard row.

**283. user (2026-06-14T03:58:13.848Z)**

The file C:/Users/avidu/Projects/bhi-wire-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**284. assistant (2026-06-14T03:58:44.273Z)**

Now a test for the additive fields + back-compat, then I'll run the full eval test set locally.

**285. assistant (2026-06-14T03:58:56.436Z)**

```
"""Tests for the honest-reporting wiring in experiment.compare (C1 + C4, additive).
```

```
Pins: compare() ADDS a `honest` block by default (CIs + paired ?F1 sign test + F1@k +
segment-count ratio) WITHOUT changing the existing fields; `honest_stats=False` restores
the legacy shape; and record_experiment carries the honest ?F1.
"""
```

```
from backend.eval.experiment import CONTROL, Variant, VideoCandidates, compare,
record_experiment
```

```
GATE = Variant(name="gateA", min_crossings=2)
```

```
def _vid(name, net_crossings):
    # 3 candidates; the MIDDLE is a false fire (net_crossings 0) absent from the golden set,
    # so Gate-A (>=2) drops it: precision/F1 up, over-segmentation 1.5 -> 1.0.
    return VideoCandidates(
        name=name,
        intervals=[(0.0, 10.0), (20.0, 30.0), (40.0, 50.0)],
        features={"net_crossings": net_crossings},
        gts=[(0.0, 10.0), (40.0, 50.0)],
    )
```

```
def _corpus():
    return [_vid("v1", [5.0, 0.0, 4.0]), _vid("v2", [3.0, 0.0, 6.0])]
```

```
def test_compare_adds_honest_block_without_touching_legacy_fields():
    res = compare(_corpus(), CONTROL, GATE)
    # legacy fields intact (no regression)
    for k in ("iou", "label_set_hash", "num_videos", "variant_a", "variant_b", "delta",
              "delta_pct", "per_video_delta"):
        assert k in res
    h = res["honest"]
    assert {"variant_a_ci", "variant_b_ci", "paired_delta_f1", "segmentation"} <= set(h)

    pd = h["paired_delta_f1"]
    assert pd["mean_delta"] > 0.0 # Gate-A removes the FP ? F1 up
    assert pd["ci_excludes_zero"] is True
    assert pd["sign_test"]["wins"] == 2.0 and pd["sign_test"]["losses"] == 0.0
```



```

# C4: CONTROL over-segments (3 preds / 2 rallies = 1.5); Gate-A is 1:1.
assert h["segmentation"]["variant_a"]["segment_count_ratio"] == 1.5
assert h["segmentation"]["variant_b"]["segment_count_ratio"] == 1.0
seg_a = h["segmentation"]["variant_a"]
assert "f1@10" in seg_a and "f1@25" in seg_a and "f1@50" in seg_a

# CI block present per metric
assert "f1" in h["variant_b_ci"] and "ci_low" in h["variant_b_ci"]["f1"]

def test_honest_stats_false_restores_legacy_shape():
    res = compare(_corpus(), CONTROL, GATE, honest_stats=False)
    assert "honest" not in res
    assert "delta" in res and "per_video_delta" in res    # legacy fields still present

def test_record_experiment_carries_honest_delta(tmp_path):
    res = compare(_corpus(), CONTROL, GATE)
    row = record_experiment(res, path=str(tmp_path / "lb.json"),
                           timestamp="2026-06-14T00:00:00Z")
    assert "honest_delta_f1" in row
    assert "sign_test" in row["honest_delta_f1"] and "ci_low" in row["honest_delta_f1"]

```

#### 286. user (2026-06-14T03:58:56.471Z)

File created successfully at: C:/Users/avidu/Projects/bhi-wire-wt/tests/test\_experiment\_honest.py (file state is current in your context ? no need to Read it back)

#### 287. user (2026-06-14T03:59:48.645Z)

```

..... [100%]
22 passed in 3.58s
=== ruff + mypy ===
All checks passed!
backend\sports\__init__.py:7: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

#### 288. assistant (2026-06-14T04:00:22.582Z)

22 tests pass (3 new + 19 existing ? no regression), ruff + mypy clean. Let me add a brief note to `EXPERIMENT\_HARNESS.md` so the doc reflects compare()'s new contract, then commit.

#### 289. user (2026-06-14T04:00:23.986Z)

```

131:### 5.1 Offline A/B (primary, label-driven, zero production risk)
132:`compare(videos, variant_A, variant_B)` ? a thin driver composing existing functions:
134:+ `calibration.leave_one_video_out()` + `metrics.score()`, graded vs golden
(`gt_loader.load_gt_intervals`).

```

#### 290. user (2026-06-14T04:00:27.852Z)

```

131    ### 5.1 Offline A/B (primary, label-driven, zero production risk)
132    `compare(videos, variant_A, variant_B)` ? a thin driver composing existing
functions:
133    `query_candidate_segments()` (load persisted features) ? re-score each variant ?
`calibration.improvement_report()`
134    + `calibration.leave_one_video_out()` + `metrics.score()`, graded vs golden
(`gt_loader.load_gt_intervals`).
135    Reports per-video and LOVO-honest aggregate IoU/P/R/F1 deltas. This is the A/B

```

test, and it needs  
136 almost no new scoring code\*\* ? only the glue.  
137  
138 ### 5.2 Shadow (for cues that touch detection, e.g. the person detector)

#### 291. user (2026-06-14T04:00:39.729Z)

The file C:/Users/avidu/Projects/bhi-wire-wt/docs/EXPERIMENT\_HARNESS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 292. assistant (2026-06-14T04:00:45.681Z)

Now commit, push, open PR, and self-merge on green CI.

#### 293. user (2026-06-14T04:00:55.701Z)

```
=== status ===
M backend/eval/experiment.py
M docs/EXPERIMENT_HARNESS.md
?? tests/test_experiment_honest.py
warning: in the working copy of 'tests/test_experiment_honest.py', LF will be replaced by CRLF
the next time Git touches it
To https://github.com/avidullu/badminton-highlight-indexer.git
* [new branch] feat/harness-honest-reporting -> feat/harness-honest-reporting
```

#### 294. user (2026-06-14T04:01:11.760Z)

<https://github.com/avidullu/badminton-highlight-indexer/pull/129>

#### 295. user (2026-06-14T04:02:59.705Z)

```
Ruff lint      pass    9s      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27487759352/job/81247121115
Syntax sweep + import smoke (no GPU stack)      pass    21s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487759352/job/81247121069
Full test suite (offline, mocked)      pass    1m31s  https://github.com/avidullu/badminton-
highlight-indexer/actions/runs/27487759352/job/81247121075
Mypy (lenient baseline) pass    35s      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27487759352/job/81247121085
Ruff lint      pass    9s      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27487759352/job/81247121115
Syntax sweep + import smoke (no GPU stack)      pass    21s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27487759352/job/81247121069
exit:0
```

#### 296. assistant (2026-06-14T04:03:03.781Z)

All green. Merging #129 and cleaning up.

#### 297. user (2026-06-14T04:03:14.035Z)

```
{"mergeStateStatus":"CLEAN","mergeable":"MERGEABLE"}
{"mergedAt":"2026-06-14T04:03:10Z","state":"MERGED"}
worktree removed
Deleted branch feat/harness-honest-reporting (was 3088022).
f680b8d feat(eval): wire C1+C4 honest reporting into experiment.compare() (additive) (#129)
8398e52 docs(archive): first end-to-end A/B experiment on the golden corpus (n=6) (#128)
d7b1f51 feat(eval): honest ?F1 ? bootstrap CI + paired sign test on the fusion lift (#127)
```

#### 298. assistant (2026-06-14T04:03:20.948Z)

Both deliverables merged: \*\*#128\*\* (archived first-A/B doc) and \*\*#129\*\* (the (b) harness wiring), master at `f680b8d`. Let me record this in the memory checkpoint.

### 299. user (2026-06-14T04:03:21.752Z)

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (FUSION P3/P4 MEASUREMENT BUILT + GPU EXTRACTION RUNNING ?
master `d7b1f51`; suite 514 green).
13     The "modeling" track for the person-cue (P3) + audio (P4) fusion lift ? all default-OFF,
no-special-casing, harness-measured.
```

### 300. user (2026-06-14T04:03:42.712Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

### 301. assistant (2026-06-14T04:04:03.316Z)

Done ? ran (a), archived it for posterity (#128), and shipped (b) (#129). Both merged to `master`.

#### (a) First golden A/B ? n=6, 100% offline (no GPU)

Re-scored the persisted WASB trajectories of all 6 golden matches, one candidate set re-scored three ways. \*\*The "persist once, re-score forever" thesis worked end-to-end.\*\*

**\*\*Per-video held-out F1:\*\***

```
video	CONTROL	Gate-A	learned
adarsh_avi	0.597	0.722	0.613
kushagra (multi-court)	0.154	0.163	0.163
largetest_doubles	0.593	0.667	0.535
gbaaddy (multi-court)	0.231	0.274	0.222
testlarge_short	0.300	0.375	0.000
mahadevpura	0.455	0.571	0.308
**mean**	**0.388**	**0.462**	**0.307**
```

**\*\*Wins / losses:\*\***

```
- ? **WIN ? Gate-A (net-crossings ? 2):** ?F1 **0.074, wins all 6 matches, sign-test p=0.031,
CI [+0.042, +0.104]** (clears 0). Over-segmentation (segment-count ratio) **1.68 ? 1.01** ?
nearly eliminated. **Promotable.**
- ? **NOT-A-WIN ? window-level 5-feature logistic:** ?F1 ?0.081, 2W/4L. `testlarge` zeroed out
(predicted no rallies) ? a **threshold/calibration artifact**, the textbook C2 motivation.
*(Distinct from the documented per-frame prototype `rally_seq_proto` at 0.583 ? a different,
stronger model; no contradiction.)*
- ? **The framework held up:** C1 gave a confident **yes** on Gate-A and a clear **no** on the
logistic ? the right call **both** directions, where a bare mean would have hidden Gate-A's
significance. C4 surfaced over-segmentation temporal-IoU buried. Multi-court is windowing-capped
(recall ceilings 0.125 / 0.312).
```

#### Posterity doc ? [#128](https://github.com/avidullu/badminton-highlight-indexer/pull/128) ? merged

`docs/archives/quality-iterations/2026-06-first-ab-experiment/` ? full method, corpus, result tables, wins/losses, lessons; + archives-index entry + a live `QUALITY\_ITERATIONS.md` pointer. The honest empirical backbone for the paper aim.

`experiment.compare()` now emits a `**honest` block by default**` (additive ? existing fields byte-unchanged; ``honest_stats=False` = legacy): per-metric bootstrap **CIs**, paired **F1 CI + exact sign test** (C1), and **F1@k + segment-count ratio** (C4); `record_experiment` carries the honest ?F1. 19 existing tests pass + 3 new; all CI green. Complements sibling **#127** (which wired the same C1 into `fusion_compare`) and reuses the #124 helpers ? no duplication.`

**\*\*Natural next steps (owner-gated):\*\*** promote Gate-A through the no-regression gate (it's a real measured win), and fix the logistic with C2 calibration + threshold-tuning-inside-LOVO. Really satisfying that the design held up under real data. ?

please do ahead with the derived recommendations, promoting Gate A with clear doc trail and rationale and fixing the logistic and threshold tuning. After this we should run the tool on "C:\Users\avidu\Projects\badminton-highlight-indexer\output\GX010125\_1080p.mp4" after isolating it into a separate directory (to avoid polluting the past work we did on this video). I believe it was 4K originally but if all the analysis is being done on the 1080p video then we can just reuse it rather than relying on the large 4K footage. I would like you to also plan a next step where we would always make a named directory for each video and put all of its outputs in it rather than splitting it in the same directory. So that a directory with 5 videos results in 5 subdirectories which contain all the info including runlogs, rallies detected, video segments, pngs etc.

This is a big, multi-part request. Let me ground it in reality first ? several parts (especially running the indexer on the video) depend on the GPU/WSL/Gemini stack, so let me probe the environment and the existing GX010125 assets before committing to a sequence.

```

=== heavy deps available here? ===
torch 2.6.0+cu124 cuda True
cv2 4.13.0
=== WSL available? ===
D:\e\faul\lt\ D:\i\s\t\r\i\b\u\t\i\o\n\:\ U\b\u\n\t\u\
█
D\d\e\f\a\u\l\l\t\ V\e\r\s\i\o\n\:\ 2█
█
█=== Gemini key in .env? ===
1
=== the target video ===
-rw-r--r-- 1 avidu 197609 2352117156 Jun 12 22:29 output/GX010125_1080p.mp4
=== existing GX010125 outputs (past work to NOT pollute) ===
output/GX010125_1080p.mp4
output/GX010125_highlights.mp4
output/GX010125_highlights_min5.mp4
output/gx010125_min5_stitch.log
output/gx010125_run.log
output/gx010125_transcode.log
output/wasb/GX010125_1080p__89605e80ed01_wasb.csv
output/wasb/GX010125_1080p__89605e80ed01_wasb_wasbcache
=== current config.json segmenter/fusion ===
default_segmenter= gemini
fusion= None

```

This is the owner's real machine with the full stack ? **\*\*torch+CUDA, cv2, WSL Ubuntu, Gemini key all present\*\*** ? and crucially **\*\*GX010125's 1080p WASB trajectory is already cached\*\***

(`output/wasb/GX010125\_1080p\_\_?\_wasb.csv` + wasbcache). So your instinct is right: we reuse the 1080p, no 4K needed, and the expensive detection is already done.

This is a large multi-stage request. Let me lay out the sequence and read the code I need for the engineering pieces and the run.

### 306. user (2026-06-14T04:12:23.775Z)

```
1     import os
2     import sys
3     import argparse
4     import logging
5     import threading
6     import traceback
7     import uuid
8     from concurrent.futures import ThreadPoolExecutor
9     import uvicorn
10    from dotenv import load_dotenv
11
12    load_dotenv()
13
14    # Centralized logging setup for the indexer (hot paths use logger, CLI output uses print
for direct user-facing messages).
15    logging.basicConfig(
16        level=logging.INFO,
17        format='%(asctime)s [%(levelname)s] %(name)s: %(message)s',
18        datefmt='%H:%M:%S'
19    )
20    logger = logging.getLogger(__name__)
21    from fastapi import FastAPI, HTTPException, Request
22    from fastapi.middleware.cors import CORSMiddleware
23    from pydantic import BaseModel
24    from typing import List, Dict, Any, Optional
25
26    from fastapi.staticfiles import StaticFiles
27    from fastapi.responses import FileResponse, HTMLResponse, JSONResponse
28    from backend.pipeline.validators.base import registry
29    from backend.infrastructure.database import Database
30    from backend.pipeline.segmenters.base import segmenter_registry
31    from backend.pipeline.stitcher.compiler import VideoCompiler
32    from backend.utils.hashing import compute_video_id # canonical file-identity hashing
33    from backend.utils.paths import resolve_under_root, safe_output_filename,
validate_input_path
34    from backend.config import load_config as load_app_config, AppConfig, StitchingConfig #
new typed config
35    from backend.results import Failure # standardized error/result contract
36    from backend.reporting import emit_indexer_report
37
38    app = FastAPI(title="Sports Highlight Indexer API")
39
40    # Enable CORS for the local web interface. allow_origins='' with allow_credentials=True
is
41    # rejected by browsers per the fetch spec and is an unsafe default to carry into the
42    # auth/tenancy work; this tool uses no cookies, so credentials stay off.
43    app.add_middleware(
44        CORSMiddleware,
45        allow_origins=["*"],
46        allow_credentials=False,
47        allow_methods=["*"],
48        allow_headers=["*"],
49    )
50
51    # Serves static frontend files directly (now under api/ per approved structure)
52    os.makedirs("backend/api/static", exist_ok=True)
53    app.mount("/static", StaticFiles(directory="backend/api/static"), name="static")
```

```

54
55     @app.get("/", response_class=HTMLResponse)
56     def serve_index():
57         index_path = "backend/api/static/index.html"
58         if os.path.exists(index_path):
59             with open(index_path, "r", encoding="utf-8") as f:
60                 return f.read()
61         return "<h3>Sports Highlight Indexer Server is Running. Please create index.html in
backend/api/static/</h3>"
62
63     # Global variables initialized at startup
64     db_instance: Optional[Database] = None
65     global_config: Optional[AppConfig] = None # now typed (Pydantic model)
66
67     # Async job executor (DEFERRED_HARDENING_PLAN #16). max_workers=1 on purpose: a single
68     # self-hosted box serializes GPU work, and the Gemini provider is rate-fragile under
69     # parallel uploads (2026-06-10 ops learning: serial uploads + backoff). The segmenters
70     # already parallelize their AI handoff internally via their own pools.

```

### 307. assistant (2026-06-14T04:12:48.252Z)

Let me understand the proven run mechanics (the owner's `scratch/run\_gx010125.py`) and the CLI's output-dir support before I sequence the run and the per-video-directory plan.

### 308. user (2026-06-14T04:12:50.846Z)

=== how the owner runs a single video E2E (the proven recipe) ===

"""One-off non-interactive E2E runner for GX010125.MP4 (modeled on run\_process\_and\_stitch.py).

Every knob ? segmenter, min\_shot\_count, paddings, watermark ? comes from config.json.

Nothing is overridden here, so the run demonstrates the config is actually picked up.

"""

```

import json
import logging
import os
import sys
import time

```

```

sys.path.insert(0, os.path.dirname(os.path.dirname(os.path.abspath(__file__)))) # repo root

```

```

from dotenv import load_dotenv

```

```

load_dotenv()

```

```

logging.basicConfig(
    level=logging.INFO,
    format="%(asctime)s [%(name)s] %(levelname)s: %(message)s",
    datefmt="%H:%M:%S",
)

```

```

from backend.infrastructure.database import Database
from backend.config import load_config
from backend.utils.hashing import compute_video_id
from backend.pipeline.segmenters import segmenter_registry
from backend.pipeline.stitcher.compiler import VideoCompiler

```

**Segmentation runs on the 1080p proxy (4K60 decode was the bottleneck ? measured frame extraction = ~2.5h for the trajectory pass). Timestamps transfer 1:1 (same timeline), so the final reel is cut from the ORIGINAL 4K for full quality.**

```

INDEX_PATH = r"C:\Users\avidu\Projects\badminton-highlight-indexer\output\GX010125_1080p.mp4"
SOURCE_PATH = r"F:\GoPro Hero Black 12\Khe\Sutra\Badminton 12 June 2026\GX010125.MP4"
OUTPUT_FILENAME = "GX010125_highlights.mp4"

```

```

def main():
    cfg = load_config()
    print(f"[CONFIG] segmenter={cfg.indexing.default_segmenter} | "
          f"min_shot_count={cfg.stitching.min_shot_count} | "
          f"begin_padding={cfg.stitching.begin_padding}s
end_padding={cfg.stitching.end_padding}s | "
          f"watermark={'on' if cfg.stitching.watermark.enabled else 'off'}", flush=True)

    for p in (INDEX_PATH, SOURCE_PATH):
        if not os.path.exists(p):
            raise SystemExit(f"video not found: {p}")

    os.makedirs(cfg.output.output_dir, exist_ok=True)
    db = Database(os.path.join(cfg.output.output_dir, cfg.output.db_name))

    t0 = time.time()
    print("[HASH] computing MD5 of the 1080p proxy...", flush=True)
    video_id = compute_video_id(INDEX_PATH)
    print(f"[HASH] video_id={video_id} ({time.time() - t0:.0f}s)", flush=True)

    segments = []
    cached = db.get_video(video_id)
    if cached and cached.get("status") == "processed":
        segments = db.query_play_segments(video_id, {})
=== main.py: output-dir / CLI run support ===
backend/main.py:227:         # report.output_dir). Never raises. Surface the paths in the
response.
backend/main.py:450:         (which only writes into output_dir; the browser previously got back an
unreachable
backend/main.py:458:         output_dir = getattr(getattr(global_config, "output", None),
"output_dir", "output") or "output"
backend/main.py:459:         path = os.path.join(output_dir, name)
backend/main.py:461:         path = resolve_under_root(path, output_dir)
backend/main.py:503:         # path too ? parity with run_process_and_stitch.py. Segments WITHOUT
shot_count
backend/main.py:529:         output_dir = getattr(getattr(global_config, "output", None),
"output_dir", "output") or "output"
backend/main.py:530:         compiler = VideoCompiler(output_dir,
begin_padding=stitching.begin_padding,
backend/main.py:539: def run_cli_diagnose_clip(video_path: str, timestamp: float):
backend/main.py:612:     parser.add_argument("--video-path", help="Local path to the MP4 sports
video file")
backend/main.py:613:     parser.add_argument("--output-dir", default="output", help="Directory
where processed clips/highlights are saved")
backend/main.py:614:     parser.add_argument("--setup", action="store_true", help="Launch
dependency checker and setup utility")
backend/main.py:615:     parser.add_argument("--debug-clip", type=float, help="Timestamp (in
seconds) of a video clip to run detailed diagnostics on")
backend/main.py:616:     parser.add_argument("--server", action="store_true", help="Run the
FastAPI local API server")
backend/main.py:617:     parser.add_argument("--port", type=int, default=8000, help="Port to run
the FastAPI server on")
backend/main.py:618:     parser.add_argument("--max-frames", type=int, help="Limit processing to
the first N sub-sampled frames for debugging")
backend/main.py:621:     parser.add_argument("--segmenter", choices=available_segmenters,
help=f"Choose rally detection segmenter. Available: {available_segmenters}")
backend/main.py:622:     parser.add_argument("--trim-start", type=float, help="Start time in
seconds for the debug trim segment feature")
backend/main.py:623:     parser.add_argument("--trim-end", type=float, help="End time in seconds
for the debug trim segment feature")
backend/main.py:624:     parser.add_argument("--trim-output", help="Output filename for the debug
trimmed segment (saves in output-dir unless absolute path)")
backend/main.py:637:         print("[ERROR] Please provide --video-path to extract a
segment.")
backend/main.py:647:         out_dir = args.output_dir

```

```

backend/main.py:669:         print("[ERROR] Please provide --video-path to debug a specific
clip.")
backend/main.py:678:     # The CLI --output-dir overrides the configured output directory. It now
lives
backend/main.py:679:     # on the typed model (OutputConfig.output_dir), so every consumer ?
including

```

### 309. assistant (2026-06-14T04:14:46.948Z)

This is the owner's GPU box and GX010125's 1080p WASB trajectory is cached ? so the run reuses 1080p, no 4K. This is a sizable multi-stage request; here's the sequence I'll execute, being honest that the video run is the long pole:

1. **\*\*Fix the logistic + threshold tuning\*\*** (C2 calibration + threshold-inside-LOVO), validated on the n=6 golden data.
2. **\*\*Promote Gate-A\*\*** ? leaderboard record + doc-trail/rationale + recommended config.
3. **\*\*Per-video output-directory plan\*\*** (design doc; I also found the relevant wart ? trajectory/handoff dirs are hardcoded to `output/`, not the configured dir).
4. **\*\*Run GX010125\_1080p\*\*** into an isolated `output/GX010125\_run/`, reusing the cached 1080p WASB.

Let me read the rest of the proven runner and the WASB cache-reuse so I set the run up correctly.

### 310. user (2026-06-14T04:14:49.286Z)

```

=== rest of run_gx010125.py ===
    segments = db.query_play_segments(video_id, {})
    if segments:
        print(f"[CACHE] reusing {len(segments)} cached segments for {video_id}", flush=True)

    if not segments:
        print("[INDEX] clean indexing (validators bypassed for the massive file, "
              "as in run_process_and_stitch.py)...", flush=True)
        db.clear_video_data(video_id)
        db.add_video(video_id, INDEX_PATH, "processed", [])
        seg_name = cfg.indexing.default_segmenter
        segmenter_cls = segmenter_registry.get(seg_name)
        if not segmenter_cls:
            raise SystemExit(f"segmenter '{seg_name}' not registered")
        segmenter = segmenter_cls(db, cfg)
        segments = segmenter.process_video(video_id, INDEX_PATH)
        print(f"[INDEX] {len(segments)} play segments indexed "
              f"(total {time.time() - t0:.0f}s).", flush=True)

    if not segments:
        raise SystemExit("no play segments detected ? nothing to stitch")

    min_shots = cfg.stitching.min_shot_count
    time_pairs, skipped = [], 0
    for s in segments:
        meta = s.get("metadata") or {}
        if isinstance(meta, str):
            try:
                meta = json.loads(meta)
            except Exception:
                meta = {}
        sc = meta.get("shot_count", 0)
        line = (f" segment id={s.get('id')} "
               f"[{float(s['start_time']):8.2f}s - {float(s['end_time']):8.2f}s] "
               f"shot_count={sc} confidence={meta.get('shot_count_confidence', '?')}")
        if sc >= min_shots:
            time_pairs.append((s["start_time"], s["end_time"]))
            print(line + " -> KEEP", flush=True)
        else:

```



```

        skipped += 1
        print(line + " -> drop", flush=True)

    print(f"[FILTER] config min_shot_count={min_shots}: "
          f"kept {len(time_pairs)}, dropped {skipped}", flush=True)
    if not time_pairs:
        raise SystemExit(f"no segments passed the min_shot_count={min_shots} filter")

    compiler = VideoCompiler(cfg.output.output_dir,
                             begin_padding=cfg.stitching.begin_padding,
                             end_padding=cfg.stitching.end_padding,
                             watermark=cfg.stitching.watermark)
    out = compiler.compile_highlights(SOURCE_PATH, time_pairs, OUTPUT_FILENAME) # cut from the
4K original
    print(f"[DONE] highlight reel: {os.path.abspath(out)} (total {time.time() - t0:.0f}s)",
          flush=True)

if __name__ == "__main__":
    main()
=== does wasb_runner reuse a cached trajectory CSV? ===
7: 4. reuses the (model-agnostic) trajectory->action-window logic from tracknet_runner.
41: keep_frames: bool = False # retain staged video + frame cache after success (debug only)
103: def run_predict(self, video_win_path: str, output_win_dir: str,
107:     All WSL/cache/output artifacts are namespaced by ``video_id`` (the file MD5) so
109:     can never reuse each other's staged frames, resumable cache, or CSV.
141:     "may fill the disk ? consider running tools/cleanup_caches.py.",
162:     if not os.path.exists(expected_csv_win):
168:     # the (huge) decoded frame cache are now pure intermediates. Delete them unless
170:     # so the frame cache + detector cache remain for a resumable re-run.
172:     logger.info("Cache hygiene: removing staged video + frame cache for %s ")

```

### 311. assistant (2026-06-14T04:15:22.832Z)

Confirmed the runner flow + that `wasb\_runner` reuses a cached trajectory CSV if present (line 162) ? critical, since a cold WASB pass is ~2.5h. Let me verify the exact reuse path so I don't accidentally trigger a full re-detection.

### 312. user (2026-06-14T04:15:24.176Z)

```

150         f"--sport {shlex.quote(self.cfg.sport)} "
151         f"--out {wsl_tilde_quote(wsl_out_dir + '/' + key + '_wasb.csv'}}"
152     )
153     logger.info("Running WASB inference on %s ...", base)
154     try:
155         res = self._wsl_bash(cmd)
156     except subprocess.TimeoutExpired:
157         logger.error("WASB inference timed out after %ss.", self.cfg.timeout_sec)
158         return None
159     if res.returncode != 0:
160         logger.error("WASB inference failed:\n%s", (res.stderr or
res.stdout)[-2000:])
161         return None
162     if not os.path.exists(expected_csv_win):
163         logger.error("WASB finished but expected CSV not found at %s",
expected_csv_win)
164         return None
165     logger.info("WASB trajectory CSV: %s", expected_csv_win)
166
167     # Success ? the trajectory CSV is the durable output; the staged video copy and
168     # the (huge) decoded frame cache are now pure intermediates. Delete them unless
169     # explicitly retained. On any earlier failure we return above WITHOUT cleaning,
170     # so the frame cache + detector cache remain for a resumable re-run.
171     if not self.cfg.keep_frames:
172         logger.info("Cache hygiene: removing staged video + frame cache for %s ")

```

```

173         "(set indexing.wasb.keep_frames=true to retain).", key)
174         self._wsl_rm_rf([frames_dir, staged_video])
175     return expected_csv_win
176
177     # Trajectory parsing + windowing are inherited from tracknet_runner (model-
agnostic).
178     parse_trajectory_csv = staticmethod(TrackNetRunner.parse_trajectory_csv)
179     trajectory_to_action_windows =
staticmethod(TrackNetRunner.trajectory_to_action_windows)

```

### 313. assistant (2026-06-14T04:15:32.561Z)

I need the start of `run\_predict` ? line 162 is a \*post\*-run check, so I must confirm there's an early cache-hit return before it triggers a 2.5h WASB pass.

### 314. user (2026-06-14T04:15:33.289Z)

```

103     def run_predict(self, video_win_path: str, output_win_dir: str,
104                     video_id: Optional[str] = None) -> Optional[str]:
105         """Run WASB on a video. Returns the Windows path to the trajectory CSV, or None.
106
107         All WSL/cache/output artifacts are namespaced by ``video_id`` (the file MD5) so
108         two videos that share a filename (e.g. two ``match.mp4`` from different folders)
109         can never reuse each other's staged frames, resumable cache, or CSV.
110         """
111         # Resolve relative dirs (the hybrids pass "output/wasb") BEFORE the /mnt
112         # translation: to_wsl_mnt_path passes relative paths through unchanged, so
113         # WSL resolved them against the inference cwd (~/.models/WASB-SBDT/src) and
114         # the CSV landed inside WSL while Windows polled the repo-relative path.
115         output_win_dir = os.path.abspath(output_win_dir)
116         os.makedirs(output_win_dir, exist_ok=True)
117         # Normalize for cross-platform basename (tests use Windows paths on Linux).
118         video_win_path = str(video_win_path).replace("\\", "/")
119         base = os.path.basename(video_win_path)
120         stem, ext = os.path.splitext(base)
121         # Unique, content-derived key: same filename + different bytes -> different key.
122         key = f"{stem}__{video_id[:12]}" if video_id else stem
123         wsl_out_dir = to_wsl_mnt_path(output_win_dir)
124         expected_csv_win = os.path.join(output_win_dir, f"{key}_wasb.csv")
125
126         if not self._stage_infer_script():
127             logger.error("Could not stage wasb_infer.py into WSL.")
128             return None
129         staged_video = self._stage_video(video_win_path, f"{key}{ext}")
130         if staged_video is None:
131             return None
132         frames_dir = f"{self.cfg.wsl_stage_dir}/{key}_frames"
133
134         # Disk guard: a 4K clip can extract ~100 GB of PNG frames. Warn early if the
135         # WSL stage filesystem is already low so the user can clean up before it fills.
136         if self.cfg.min_free_gb > 0:
137             free = self._wsl_free_gb(self.cfg.wsl_stage_dir)
138             if free is not None and free < self.cfg.min_free_gb:
139                 logger.warning(
140                     "Low WSL disk: %.1f GB free at %s (< min_free_gb=%.1f). Frame
extraction "
141                     "may fill the disk ? consider running tools/cleanup_caches.py.",
142                     free, self.cfg.wsl_stage_dir, self.cfg.min_free_gb)
143
144         cmd = (
145             f"conda activate {self.cfg.conda_env} && cd {self.cfg.repo_dir}/src && "
146             f"python wasb_infer.py "
147             f"--video {wsl_tilde_quote(staged_video)} "
148             f"--frames_out_dir {wsl_tilde_quote(frames_dir)} "
149             f"--weights {wsl_tilde_quote(self.cfg.weights_path)} "

```

150 f"--sport {shlex.quote(self.cfg.sport)} "

### 315. assistant (2026-06-14T04:17:03.365Z)

Critical finding before I run anything: `wasb\_runner.run\_predict` has **no** early cache-hit return ? the standard pipeline would **re-run** WASB from scratch (~tens of min on 1080p), overwriting the cached trajectory (i.e., polluting it). The cached CSV is only directly reusable via **offline re-scoring**. So the honest "reuse the 1080p, don't pollute" path for the run is to re-score the cached trajectory offline.

Given the scope, let me sequence pragmatically. I'll start with the **logistic fix** (the core engineering, validatable on golden now), in a worktree. Let me set it up and implement the additive C2-calibration + threshold-in-LOVO.

### 316. user (2026-06-14T04:17:08.907Z)

```
f680b8d feat(eval): wire C1+C4 honest reporting into experiment.compare() (additive) (#129)
HEAD is now at f680b8d feat(eval): wire C1+C4 honest reporting into experiment.compare()
(additive) (#129)
=== current learned path: predict() + evaluate_variant threshold handling ===
102: - ``threshold`` ? classifier probability cutoff. ``merge_gap`` ? post-filter
114:   threshold: float = 0.5
131:   "threshold": self.threshold,
205: def predict(variant: Variant, vc: VideoCandidates,
223:     proba = model.predict_proba(_feature_matrix(vc, variant.feature_names))
224:     mask = [m and (float(proba[i]) >= variant.threshold) for i, m in enumerate(mask)]
232: def _train(variant: Variant, videos: Sequence[VideoCandidates], iou: float) ->
LogisticRegression:
269:   recipe_lovo = bool(variant.feature_names) and variant.classifier_model is None
273:   if recipe_lovo and not honest:
274:     all_model = _train(variant, videos, iou)
277:   if recipe_lovo and honest:
281:     model: Optional[LogisticRegression] = _train(variant, others, iou)
282:     elif recipe_lovo: # honest=False ? one model over all videos
298:     "honest_lovo": recipe_lovo and honest,
```

### 317. user (2026-06-14T04:17:18.004Z)

```
205 def predict(variant: Variant, vc: VideoCandidates,
206     model: Optional[LogisticRegression] = None) -> List[Interval]:
207     """Variant prediction: gate by persisted features, optionally filter by a learned
208     model, then merge. Pure and deterministic.
209
210     ``model`` (an already-fitted `LogisticRegression`) is supplied by `compare` per LOVO
211     fold; if omitted and the variant pins ``classifier_model``, that frozen model is
212     loaded. A learned variant with neither a supplied nor a pinned model raises ? there
213     is nothing to score with.
214     """
215     mask = _gate_mask(variant, vc)
216     if variant.feature_names:
217         if model is None:
218             if variant.classifier_model is None:
219                 raise ExperimentError(
220                     f"variant {variant.name!r} is learned but has no model to predict "
221                     f"with (pass a fitted model or set classifier_model)")
222                 model = LogisticRegression.load(variant.classifier_model)
223                 proba = model.predict_proba(_feature_matrix(vc, variant.feature_names))
224                 mask = [m and (float(proba[i]) >= variant.threshold) for i, m in
enumerate(mask)]
225                 kept = [iv for iv, keep in zip(vc.intervals, mask) if keep]
226                 return merge_overlaps(kept, gap_tol=variant.merge_gap)
227
228
229     # ----- variant scoring
```

```

230
231
232     def _train(variant: Variant, videos: Sequence[VideoCandidates], iou: float) ->
LogisticRegression:
233         """Fit the variant's classifier on the pooled candidates of ``videos`` (labels via
the
234         evaluator's own IoU matcher). The honest-LOVO training step from `distill_local`. """
235         X: List[List[float]] = []
236         y: List[int] = []
237         for vc in videos:
238             if vc.gts is None:
239                 raise ExperimentError(f"cannot train: {vc.name} has no golden labels")
240             X.extend(_feature_matrix(vc, variant.feature_names or []))
241             y.extend(label_candidates(vc.intervals, vc.gts, iou))
242         if not X:
243             raise ExperimentError("cannot train a learned variant on zero candidates")
244         return LogisticRegression().fit(X, y, variant.feature_names)
245
246
247     def evaluate_variant(videos: Sequence[VideoCandidates], variant: Variant,
248                         iou: float = 0.5, honest: bool = True) -> Dict[str, Any]:
249         """Score a variant on a labeled corpus. Returns per-video Scores + predictions + a
250         mean aggregate.
251
252         Prediction policy:
253         - **static variant** (no learned filter): predict each video directly ? no
training,
254         so no leakage; per-video scoring is already honest.
255         - **learned, pinned model** (`classifier_model` set): score every video with the
256         SHIPPED model. ? optimistic if those videos were in its training set ? use this
to
257         report "how the shipped artifact does here", not for the promotion decision.
258         - **learned recipe** (`feature_names`, no pinned model) + `honest=True`: LOVO
?
259         train on the OTHER videos, predict the held-out one. The number to trust.
260         - learned recipe + `honest=False`: train on all, predict each (overfit; sanity
only).
261         """
262         for vc in videos:
263             if vc.gts is None:
264                 raise ExperimentError(f"{vc.name} has no golden labels; use
compare_unlabeled")
265
266         per_video: List[Dict[str, Any]] = []
267         predictions: Dict[str, List[Interval]] = {}
268
269         recipe_lovo = bool(variant.feature_names) and variant.classifier_model is None
270         pinned_model = (LogisticRegression.load(variant.classifier_model)
271                         if variant.classifier_model is not None else None)
272         all_model = None
273         if recipe_lovo and not honest:
274             all_model = _train(variant, videos, iou)
275
276         for i, vc in enumerate(videos):
277             if recipe_lovo and honest:
278                 others = [v for j, v in enumerate(videos) if j != i]
279                 if not others:
280                     raise ExperimentError("LOVO needs >=2 videos; pass honest=False or a
pinned model")
281                 model: Optional[LogisticRegression] = _train(variant, others, iou)
282                 elif recipe_lovo: # honest=False ? one model over all
videos
283                     model = all_model
284                 else: # static variant or pinned model
285                     model = pinned_model

```

```

286         preds = predict(variant, vc, model=model)
287         assert vc.gts is not None # guarded at function top
288         sc = score(preds, vc.gts, iou)
289         per_video.append({"video": vc.name, "num_pred": len(preds), **{m: getattr(sc, m)
for m in _METRICS}})
290         predictions[vc.name] = preds
291
292     aggregate = {m: (sum(p[m] for p in per_video) / len(per_video) if per_video else
0.0)
293                 for m in _METRICS}
294     return {
295         "variant": variant.name,
296         "spec_hash": variant.spec_hash(),
297         "is_learned": variant.is_learned(),
298         "honest_lovo": recipe_lovo and honest,
299         "per_video": per_video,
300         "aggregate": aggregate,

```

### 318. assistant (2026-06-14T04:17:17.128Z)

Let me read the exact `predict()` + `evaluate\_variant` code to edit precisely.

### 319. assistant (2026-06-14T04:17:49.973Z)

Now I'll implement the fix additively (defaults preserve current behavior ? no-regression): a `tune\_threshold` (operating point chosen on the \*training\* fold) + `calibrate` (Platt/isotonic) on the learned LOVO path.

### 320. user (2026-06-14T04:17:54.982Z)

The file C:/Users/avidu/Projects/bhi-fix-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

### 321. user (2026-06-14T04:17:58.601Z)

```

104         overlap-merge gap (seconds), the same `gemini_refine.merge_overlaps` default.
105     """
106
107     name: str
108     segmenter: str = "wasb_hybrid"
109     config_overrides: Dict[str, Any] = field(default_factory=dict)
110     cue_flags: Dict[str, bool] = field(default_factory=dict)
111     min_crossings: int = 0
112     min_players: int = 0
113     classifier_model: Optional[str] = None
114     feature_names: Optional[List[str]] = None
115     threshold: float = 0.5
116     merge_gap: float = 1.0
117
118     def is_learned(self) -> bool:
119         """True if this variant applies a learned classifier (pinned model or
recipe)."""
120         return self.classifier_model is not None or bool(self.feature_names)
121
122     def to_dict(self) -> dict:
123         return {
124             "name": self.name,
125             "segmenter": self.segmenter,
126             "config_overrides": self.config_overrides,
127             "cue_flags": self.cue_flags,
128             "min_crossings": self.min_crossings,
129             "min_players": self.min_players,
130             "classifier_model": self.classifier_model,
131             "feature_names": self.feature_names,

```

```
132         "threshold": self.threshold,
133         "merge_gap": self.merge_gap,
134     }
135
136     @classmethod
137     def from_dict(cls, d: dict) -> "Variant":
```

322. user (2026-06-14T04:18:06.364Z)

The file C:/Users/avidu/Projects/bhi-fix-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

323. user (2026-06-14T04:18:11.454Z)

The file C:/Users/avidu/Projects/bhi-fix-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

324. assistant (2026-06-14T04:18:14.941Z)

Now `predict()` (threshold/calibrator overrides) + the train-fold calibration/threshold helper + the LOVO loop:

325. user (2026-06-14T04:18:30.937Z)

The file C:/Users/avidu/Projects/bhi-fix-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

326. user (2026-06-14T04:18:38.991Z)

The file C:/Users/avidu/Projects/bhi-fix-wt/backend/eval/experiment.py has been updated successfully. (file state is current in your context ? no need to Read it back)

327. assistant (2026-06-14T04:18:58.213Z)

Now validate the fix on the n=6 golden data ? does `tune\_threshold`/`calibrate` stop the learned variant from zeroing out (testlarge) and improve the mean? Running against the worktree's new code.

328. user (2026-06-14T04:19:14.462Z)

| variant                   | adarsh_avi | kushagra_s | largetest_ | gbaaddy | testlarge_ | mahadevpur |
|---------------------------|------------|------------|------------|---------|------------|------------|
| MEAN                      |            |            |            |         |            |            |
| learned (fixed 0.5)       | 0.613      | 0.163      | 0.535      | 0.222   | 0.000      | 0.308      |
| 0.307                     |            |            |            |         |            |            |
| learned + tune_threshold  | 0.632      | 0.163      | 0.672      | 0.269   | 0.300      | 0.500      |
| 0.423                     |            |            |            |         |            |            |
| learned + tune + platt    | 0.622      | 0.160      | 0.626      | 0.264   | 0.300      | 0.500      |
| 0.412                     |            |            |            |         |            |            |
| learned + tune + isotonic | 0.632      | 0.163      | 0.626      | 0.255   | 0.300      | 0.500      |
| 0.413                     |            |            |            |         |            |            |

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

329. assistant (2026-06-14T04:19:50.479Z)

```
**The fix works and is validated on real golden data:**

variant	testlarge	MEAN F1
learned (fixed 0.5)	**0.000** (the bug)	0.307
**learned + tune_threshold**	**0.300**	**0.423**
+ platt / + isotonic	0.300	0.412 / 0.413
```

Threshold-in-LOVO un-breaks the zeroing and lifts mean F1 \*\*0.307 ? 0.423 (+0.116)\*\* ? the

learned variant now **\*\*beats CONTROL (0.388)\*\*** too. Let me add a focused test and run all gates.

### 330. assistant (2026-06-14T04:19:59.146Z)

```
"""Tests for the learned-variant fix: train-fold threshold tuning + calibration (additive).

Pins the additive/no-regression contract (defaults OFF) + the round-trip + that a tuned
variant runs and that a lower threshold keeps at least as many windows. The real proof
that it fixes the first-A/B zeroing is the golden LOVO validation (see the PR).
"""
from backend.eval.experiment import Variant, VideoCandidates, _train, evaluate_variant, predict

SH = ["duration", "peak_velocity", "mean_velocity", "active_ratio", "net_crossings"]

def _vid(name):
    return VideoCandidates(
        name=name,
        intervals=[(0.0, 10.0), (20.0, 30.0), (40.0, 50.0)],
        features={
            "duration": [10.0, 10.0, 10.0],
            "peak_velocity": [0.9, 0.1, 0.8],
            "mean_velocity": [0.5, 0.05, 0.5],
            "active_ratio": [0.8, 0.1, 0.7],
            "net_crossings": [5.0, 0.0, 4.0],
        },
        gts=[(0.0, 10.0), (40.0, 50.0)],
    )

def test_defaults_are_off_additive():
    v = Variant(name="x")
    assert v.tune_threshold is False and v.calibrate is None

def test_variant_tuning_roundtrips():
    v = Variant(name="v", feature_names=SH, tune_threshold=True, calibrate="isotonic")
    d = v.to_dict()
    assert d["tune_threshold"] is True and d["calibrate"] == "isotonic"
    assert Variant.from_dict(d).tune_threshold is True

def test_tuned_and_calibrated_variant_runs_lovo():
    vids = [_vid("a"), _vid("b")]
    for var in (Variant(name="Lt", feature_names=SH, tune_threshold=True),
                Variant(name="Ltp", feature_names=SH, tune_threshold=True, calibrate="platt"),
                Variant(name="Lti", feature_names=SH, tune_threshold=True,
                        calibrate="isotonic")):
        ev = evaluate_variant(vids, var, 0.5, honest=True)
        assert 0.0 <= ev["aggregate"]["f1"] <= 1.0

def test_predict_threshold_override_monotone():
    vc = _vid("a")
    model = _train(Variant(name="m", feature_names=SH), [vc], 0.5)
    hi = predict(Variant(name="hi", feature_names=SH), vc, model=model, threshold=0.99)
    lo = predict(Variant(name="lo", feature_names=SH), vc, model=model, threshold=0.01)
    assert len(lo) >= len(hi)          # a lower cutoff keeps at least as many windows
```

### 331. user (2026-06-14T04:19:59.160Z)

File created successfully at: C:/Users/avidu/Projects/bhi-fix-wt/tests/test\_variant\_tuning.py  
(file state is current in your context ? no need to Read it back)

### 332. user (2026-06-14T04:20:28.226Z)

```
..... [100%]
26 passed in 2.83s
=== ruff + mypy ===
All checks passed!
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\eval\experiment.py:334: error: Argument 2 to "_calibrate_and_tune" has incompatible type
"LogisticRegression | None"; expected "LogisticRegression" [arg-type]
Found 1 error in 1 file (checked 1 source file)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

### 333. assistant (2026-06-14T04:20:40.392Z)

One mypy narrowing issue ? `model` is typed `Optional`. Let me add the narrowing assert  
(`\_train` never returns None).

### 334. user (2026-06-14T04:20:42.653Z)

The file C:/Users/avidu/Projects/bhi-fix-wt/backend/eval/experiment.py has been updated  
successfully. (file state is current in your context ? no need to Read it back)

### 335. user (2026-06-14T04:21:02.139Z)

```
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 1 source file
=== commit + push ===
warning: in the working copy of 'tests/test_variant_tuning.py', LF will be replaced by CRLF the
next time Git touches it
On branch feat/learned-variant-fix
Your branch is ahead of 'origin/master' by 1 commit.
(use "git push" to publish your local commits)

nothing to commit, working tree clean
To https://github.com/avidullu/badminton-highlight-indexer.git
 * [new branch]      feat/learned-variant-fix -> feat/learned-variant-fix
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

### 336. assistant (2026-06-14T04:21:10.760Z)

Let me verify the commit landed cleanly, then open the PR.

### 337. user (2026-06-14T04:21:22.512Z)

```
cdffb9f @placeholder@
backend/eval/experiment.py | 66 ++++++-----
tests/test_variant_tuning.py | 53 ++++++
2 files changed, 113 insertions(+), 6 deletions(-)
https://github.com/avidullu/badminton-highlight-indexer/pull/130
```

### 338. assistant (2026-06-14T04:21:35.006Z)

The commit message leaked a placeholder (`@placeholder@`) due to a shell-quoting slip. Let me  
fix it with an amend + force-push to my feature branch (PR #130 already open).

### 339. user (2026-06-14T04:21:44.626Z)

```
To https://github.com/avidullu/badminton-highlight-indexer.git
+ cdffb9f...40bcc75 feat/learned-variant-fix -> feat/learned-variant-fix (forced update)
=== fixed message ===
40bcc75 feat(eval): fix the learned variant ? train-fold threshold tuning + calibration (C2)
```



#### 340. assistant (2026-06-14T04:21:51.107Z)

Message fixed. Watching CI on #130, then self-merge.

#### 341. user (2026-06-14T04:21:53.230Z)

no checks reported on the 'feat/learned-variant-fix' branch  
exit:0

#### 342. assistant (2026-06-14T04:22:01.398Z)

CI hasn't registered yet after the force-push. Let me wait briefly and re-watch.

#### 343. user (2026-06-14T04:23:47.863Z)

```
attempt 1: 4 checks registered
Syntax sweep + import smoke (no GPU stack)      pass    20s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488158457/job/81248197153
Full test suite (offline, mocked)               pass    1m29s  https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488158457/job/81248197172
Mypy (lenient baseline) pass                    32s      https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488158457/job/81248197143
Ruff lint                                       pass     13s      https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488158457/job/81248197176
Syntax sweep + import smoke (no GPU stack)      pass    20s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488158457/job/81248197153
exit:0
```

#### 344. assistant (2026-06-14T04:23:51.593Z)

All four lanes green. Merging #130 and cleaning up.

#### 345. user (2026-06-14T04:23:57.081Z)

```
{"mergedAt":"2026-06-14T04:23:56Z","state":"MERGED"}
worktree removed
Deleted branch feat/learned-variant-fix (was 40bcc75).
```

#### 346. assistant (2026-06-14T04:24:12.242Z)

#130 merged ? the learned variant is fixed and validated. Now the GX010125 run. Let me probe the 1080p's fps and re-score its **cached** trajectory into an isolated dir (reuses 1080p, no GPU re-detection, no pollution).

#### 347. user (2026-06-14T04:24:15.817Z)

```
fps=59.940 w=1920 h=1080 frames=46080
=== cached trajectory rows ===
46081
```

#### 348. assistant (2026-06-14T04:24:32.548Z)

1080p: 59.94 fps, 1920x1080, ~12.8 min, full cached trajectory. Now the isolated re-score (reuses the cached 1080p trajectory; no labels for this footage, so it's the unlabeled SxS mode ? CONTROL vs Gate-A divergence).

#### 349. assistant (2026-06-14T04:24:47.951Z)

```
""Run the framework on GX010125_1080p by REUSING its cached WASB trajectory ? offline,
no GPU, no 4K, isolated outputs (output/GX010125_run/), zero pollution of past work.
```

This is the "persist detection once, re-score forever" path on a REAL, unlabeled match:  
parse the cached 1080p trajectory -> candidate rally windows -> CONTROL vs Gate-A

```

(net-crossings >= 2, the measured precision win) side-by-side (no golden labels exist for
this footage, so we report the divergence, not F1).
"""
import json
import os

from backend.eval.distill_local import build_candidates, FEATURE_NAMES
from backend.eval.experiment import CONTROL, Variant, VideoCandidates, compare_unlabeled,
predict

TRAJ = "output/wasb/GX010125_1080p__89605e80ed01_wasb.csv"
FPS, FRAME_W = 59.94, 1920.0
OUT_DIR = "output/GX010125_run"

def _dur(ivs):
    return sum(e - s for s, e in ivs)

def _write_csv(path, ivs):
    with open(path, "w", encoding="utf-8") as f:
        f.write("start_sec,end_sec,duration_sec\n")
        for s, e in ivs:
            f.write(f"{s:.3f},{e:.3f},{e - s:.3f}\n")

def main():
    os.makedirs(OUT_DIR, exist_ok=True)
    intervals, X = build_candidates(TRAJ, FPS, FRAME_W)
    feats = {FEATURE_NAMES[k]: [r[k] for r in X] for k in range(len(FEATURE_NAMES))}
    vc = VideoCandidates(name="GX010125_1080p", intervals=intervals, features=feats, gts=None)

    gateA = Variant(name="gateA", min_crossings=2)
    ctrl_preds = predict(CONTROL, vc)
    gate_preds = predict(gateA, vc)
    sxs = compare_unlabeled(vc, CONTROL, gateA) # agreed / a_only (dropped by Gate-A) / b_only

    print("=" * 70)
    print(f"GX010125_1080p ? offline re-score of the CACHED trajectory (no GPU)")
    print("=" * 70)
    print(f" candidate windows from trajectory : {len(intervals)}")
    print(f" CONTROL rallies (merge all) : {len(ctrl_preds)} ({_dur(ctrl_preds) /
60:.1f} min)")
    print(f" Gate-A rallies (net_crossings>=2) : {len(gate_preds)} ({_dur(gate_preds) /
60:.1f} min)")
    print(f" Gate-A dropped (a_only) : {len(sxs['a_only'])} windows "
f"({_dur(sxs['a_only']) / 60:.1f} min of likely held-shuttle/feed FPs)")
    print(f" both agree (kept) : {sxs['num_agreed']}")

    _write_csv(os.path.join(OUT_DIR, "GX010125_control_rallies.csv"), ctrl_preds)
    _write_csv(os.path.join(OUT_DIR, "GX010125_gateA_rallies.csv"), gate_preds)
    summary = {
        "video": "GX010125_1080p", "source": "cached WASB trajectory (offline re-score, no
GPU)",
        "fps": FPS, "num_candidates": len(intervals),
        "control": {"num_rallies": len(ctrl_preds), "total_min": round(_dur(ctrl_preds) / 60,
2)},
        "gateA": {"num_rallies": len(gate_preds), "total_min": round(_dur(gate_preds) / 60, 2),
        "dropped_windows": len(sxs["a_only"])},
    }
    json.dump(summary, open(os.path.join(OUT_DIR, "GX010125_rescore_summary.json"), "w"),
indent=2)
    print(f"\n wrote isolated outputs -> {OUT_DIR}/ (control + gateA rally CSVs + summary)")
    print(" NOTE: no golden labels for this footage -> divergence reported, not F1. A full")
    print(" Gemini-confirmed highlight reel is a separate ~tens-of-min GPU+API run.")

```

```
if __name__ == "__main__":
    main()
```

### 350. user (2026-06-14T04:24:47.965Z)

File created successfully at: C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/run\_gx010125\_rescore.py (file state is current in your context ? no need to Read it back)

### 351. user (2026-06-14T04:24:57.795Z)

```
=====
GX010125_1080p ? offline re-score of the CACHED trajectory (no GPU)
=====
candidate windows from trajectory : 45
CONTROL rallies (merge all)       : 45 (11.0 min)
Gate-A rallies (net_crossings>=2) : 27 (9.3 min)
Gate-A dropped (a_only)           : 18 windows (1.7 min of likely held-shuttle/feed FPs)
both agree (kept)                 : 27

wrote isolated outputs -> output/GX010125_run/ (control + gateA rally CSVs + summary)
NOTE: no golden labels for this footage -> divergence reported, not F1. A full
Gemini-confirmed highlight reel is a separate ~tens-of-min GPU+API run.
=== isolated dir contents ===
total 33
drwxr-xr-x 1 avidu 197609    0 Jun 14 09:54 .
drwxr-xr-x 1 avidu 197609    0 Jun 14 09:54 ..
-rw-r--r-- 1 avidu 197609 1082 Jun 14 09:54 GX010125_control_rallies.csv
-rw-r--r-- 1 avidu 197609  672 Jun 14 09:54 GX010125_gateA_rallies.csv
-rw-r--r-- 1 avidu 197609  306 Jun 14 09:54 GX010125_rescore_summary.json
```

### 352. assistant (2026-06-14T04:25:31.946Z)

The tool ran on the real **GX010125\_1080p** (reusing the cached 1080p trajectory ? no GPU, no 4K) into an isolated `output/GX010125_run/`, no pollution: **45** candidate windows ? CONTROL 45 rallies (11.0 min) vs Gate-A 27 rallies (9.3 min); Gate-A dropped 18 low-crossing windows (~1.7 min of likely held-shuttle/feed FPs). **Now the doc trail (promote Gate-A) + the per-video-layout plan, as one docs PR.**

### 353. user (2026-06-14T04:25:36.287Z)

2e5892e feat(eval): fix the learned variant ? train-fold threshold tuning + calibration (C2) (#130)  
HEAD is now at 2e5892e feat(eval): fix the learned variant ? train-fold threshold tuning + calibration (C2) (#130)

### 354. user (2026-06-14T04:26:03.556Z)

File created successfully at: C:/Users/avidu/Projects/bhi-docs2-wt/docs/PER\_VIDEO\_OUTPUT\_LAYOUT.md (file state is current in your context ? no need to Read it back)

### 355. assistant (2026-06-14T04:26:03.525Z)

## Per-video output layout ? one directory per video (design / plan)

**\*\*Status:\*\*** **\*\*DESIGN** ? owner-gated, NOT started.**\*\*** Proposed convention so processing N videos yields N self-contained subdirectories instead of one tangled `output/`. Docs-only.

### The problem

Today every run writes into a **flat** `output/` with filename prefixes ? highlights (`<video>_highlights.mp4`), run logs (`<video>_run.log`), eval CSVs (`<video>_vs_human.csv`), the shared SQLite DB (`sports_indexer.db`), plus **hardcoded shared subdirs** the CLI

```

`--output-dir` does not redirect:
- `output/wasb/` and `output/tracknet/` ? trajectory CSVs + frame caches
  (`trajectory_hybrid.py:~176`, `out_dir = os.path.join("output", self.family)` ? literal
  `"output"`).
- `output/chunks_<provider>_handoff/` ? the AI-handoff clip staging (same file).

```

So 5 videos produce an interleaved pile, and `--output-dir` only relocates *some* artifacts (highlights/clips/DB) while trajectories and handoff clips still land in the shared `output/`. The first isolated run (`output/GX010125\_run/`, 2026-06-14) had to be done by hand and is only *partially* isolated for this reason.

## The target

```

...
output/
  <video_slug>/                # slug = sanitized stem + short video_id (collision-safe)
    runlog.txt                  # the per-video run log (today: output/<video>_run.log)
    sports_indexer.db           # per-video DB (or a row-scoped view of a shared DB ? see below)
    wasb/<key>_wasb.csv          # the trajectory (today: output/wasb/)
    handoff/                    # AI-handoff clips (today: output/chunks_*_handoff/, auto-
cleaned)
    rallies.csv                 # detected rally [start,end] + shot_count
    segments/                   # per-segment clips / pngs / debug frames
    highlights.mp4              # the reel
    report.json / report.md     # the observability report
...

```

A directory with 5 videos ? 5 subdirectories, each holding **everything** for that video.

## Design

```

- video_slug(video_path, video_id) -> str` ? a new helper (next to `utils/hashing.py`):
  sanitized stem + `video_id[:12]` (the MD5 already used as the cache key), e.g.
  `GX010125_1080p__89605e80ed01`. Collision-safe (same stem / different bytes ? different dir).
- Thread one `run_dir` through the pipeline instead of the literal `"output"`:
  - `TrajectoryHybridSegmenter.process_video` ? `out_dir = os.path.join(run_dir, self.family)`
    and the handoff dir under `run_dir` (replaces the two hardcoded `"output"` joins).
  - The `wasb_runner`/`tracknet_runner` already take `output_win_dir` ? just pass the per-video
one.
  - The CLI `--output-dir` (already in `backend/main.py`) becomes the parent; the pipeline
    creates `<output-dir>/<video_slug>/` and writes everything under it.
- DB choice (owner decision): (a) a per-video `sports_indexer.db` in the slug dir (simplest,
  fully isolated, but loses cross-video queries), or (b) keep the shared DB and only isolate
  files (the slug dir holds artifacts; the DB stays one queryable store keyed by `video_id`).
Recommend (b) ? the DB is already `video_id`-scoped; isolate the heavy file artifacts,
  keep one queryable metadata store. The slug dir then omits its own DB.
- Back-compat: default parent stays `output/`; a flag (or just the default-on behaviour)
  switches flat?per-video. Old flat outputs are left untouched; new runs go to slug dirs.

```

## Phasing (owner-gated; ONE PR per slice, off `master`)

```

- L0 ? this plan.
- L1 ? `video_slug` helper + thread `run_dir` through `trajectory_hybrid` (replace the two
  hardcoded `"output"` joins) + the per-video subdir creation; CLI `--output-dir` = parent.
  Pure plumbing; a regression test pinning that all artifacts land under the slug dir.
- L2 ? runlog + report + rallies.csv into the slug dir; migrate the eval/CSV writers.
- L3 ? cache-hygiene + cleanup_caches awareness of the new layout (per-video purge).

```

## Relation to the first isolated run

`output/GX010125\_run/` (the 2026-06-14 offline re-score) is the **manual prototype** of this layout ? it isolated the rally CSVs + summary, but the trajectory still lives in the shared `output/wasb/` precisely because of the hardcoded path L1 fixes. This plan makes that isolation automatic and complete for every video.

## 356. assistant (2026-06-14T04:26:10.421Z)

Now the Gate-A promotion record in `QUALITY\_ITERATIONS.md` + indexing the new plan in

`docs/README.md`.

### 357. user (2026-06-14T04:26:11.159Z)

```
40
41     ## Current track (live) ? newest first
42
43     ### First end-to-end A/B on the golden corpus (n=6) ? **ARCHIVED** *(2026-06-14)*
```

### 358. user (2026-06-14T04:26:32.237Z)

The file C:/Users/avidu/Projects/bhi-docs2-wt/docs/QUALITY\_ITERATIONS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

### 359. user (2026-06-14T04:26:38.242Z)

```
11     - [DEFERRED_HARDENING_PLAN.md](DEFERRED_HARDENING_PLAN.md) ? Design (pending approval)
for the larger deferred items: packaging, train/eval guardrails, and the async job model (#16,
the first infra slice toward the platform).
12     - [I18N_PLAN.md](I18N_PLAN.md) ? Design (owner-gated) for internationalizing the product
**by design**: Kannada-first, Hindi next, extensible to other Indian languages. Shared
framework-agnostic locale foundation across UI chrome, coach-facing backend text, and the future
NL-query / coach-mode surfaces.
13     - [BACKLOG.md](BACKLOG.md) ? Deferred items with rationale and revisit triggers.
```

### 360. user (2026-06-14T04:26:44.745Z)

The file C:/Users/avidu/Projects/bhi-docs2-wt/docs/README.md has been updated successfully. (file state is current in your context ? no need to Read it back)

### 361. assistant (2026-06-14T04:26:49.585Z)

Committing the docs PR (Gate-A promotion record + per-video layout plan), then self-merge on green.

### 362. user (2026-06-14T04:27:06.989Z)

warning: in the working copy of 'docs/PER\_VIDEO\_OUTPUT\_LAYOUT.md', LF will be replaced by CRLF the next time Git touches it

```
A docs/PER_VIDEO_OUTPUT_LAYOUT.md
M docs/QUALITY_ITERATIONS.md
M docs/README.md
* [new branch] docs/gate-a-promotion-and-layout -> docs/gate-a-promotion-and-layout
https://github.com/avidullu/badminton-highlight-indexer/pull/131
```

### 363. assistant (2026-06-14T04:27:11.146Z)

PR [#131](https://github.com/avidullu/badminton-highlight-indexer/pull/131). Watching CI, then self-merge + cleanup.

### 364. user (2026-06-14T04:29:18.008Z)

```
Syntax sweep + import smoke (no GPU stack)    pass    20s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488260451/job/81248456517
Full test suite (offline, mocked)              pass    1m42s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488260451/job/81248456534
Mypy (lenient baseline) pass    32s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488260451/job/81248456535
Ruff lint pass    9s https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488260451/job/81248456518
Syntax sweep + import smoke (no GPU stack)    pass    20s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27488260451/job/81248456517
```

### 365. user (2026-06-14T04:29:27.478Z)

MERGED  
worktree removed  
Deleted branch docs/gate-a-promotion-and-layout (was ab82ba3).

### 366. user (2026-06-14T04:29:36.017Z)

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (FIRST GOLDEN A/B EXPERIMENT + HARNESS HONEST-REPORTING WIRED
? **PR #128 + #129 MERGED**
```

### 367. assistant (2026-06-14T04:29:35.100Z)

All four items merged. Let me record the checkpoint, then summarize.

### 368. user (2026-06-14T04:29:55.764Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

### 369. assistant (2026-06-14T04:30:14.514Z)

All four items done and merged. Summary:

## 1 & 2 ? Derived recommendations (acted on the first A/B)

**\*\*Fixed the learned variant ? [#130](https://github.com/avidullu/badminton-highlight-indexer/pull/130) ?\*\*** (validated on n=6 golden, offline):  
- Added additive, default-OFF `Variant.tune\_threshold` (operating point picked on the **\*\*training\*\*** fold) + `calibrate` (Platt/isotonic, C2).  
- **\*\*Mean held-out F1 0.307 ? 0.423 (+0.116)\*\***, `testlarge` **\*\*0.000 ? 0.300\*\*** (un-zeroed), and the learned variant now **\*\*beats CONTROL (0.388)\*\***. The fixed-0.5 cutoff was the bug, exactly as the experiment diagnosed.

**\*\*Promoted Gate-A ? [#131](https://github.com/avidullu/badminton-highlight-indexer/pull/131) ?\*\*** with a full doc trail: the measured win (?F1 +0.074, 6/6 matches, p=0.031, over-seg 1.68?1.01), the recommended config (`indexing.fusion.min\_crossings=2`), and the **\*\*honest boundary\*\*** ? I recorded the promotion but did **\*\*not\*\*** silently flip the live served default, because the offline win is on the proposal candidate set and the production fusion path is still unverified on real video. That last flip is owner-gated pending a real fusion-path run.

## 3 ? Ran on GX010125\_1080p (isolated, reuses 1080p) ?

Key discovery: `wasb\_runner.run\_predict` has **\*\*no cache-hit early-return\*\*** ? the full pipeline would **\*\*re-run\*\*** WASB (~tens of min) and overwrite the cached trajectory (pollution). So the honest "reuse the 1080p" path is an **\*\*offline re-score\*\*** of the cached trajectory:  
- **\*\*45 candidate windows ? CONTROL 45 rallies (11.0 min) vs Gate-A 27 rallies (9.3 min)\*\***;  
Gate-A dropped **\*\*18 low-crossing windows (~1.7 min)\*\*** of likely held-shuttle/feed FPs.  
- Isolated in **\*\*`output/GX010125\_run/`\*\*** (rally CSVs + summary) ? past GX010125 work untouched. No golden labels for this footage, so I report divergence, not F1.  
- A full **\*\*Gemini-confirmed highlight reel\*\*** is a separate ~tens-of-min GPU+API run (overwrites

the trajectory cache) ? yours to kick off when ready; I can wire it to the isolated layout first.

#### 4 ? Per-video output directory plan ? [#131](https://github.com/avidullu/badminton-highlight-indexer/pull/131)

New `docs/PER\_VIDEO\_OUTPUT\_LAYOUT.md`: `output/<video\_slug>/` holding runlog, trajectory, rallies, segments, pngs, highlights per video. I found the blocker ? `trajectory\_hybrid.py` hardcodes `"output"` for the WASB/handoff dirs (so `--output-dir` doesn't fully isolate today); the plan's L1 phase fixes exactly that.

Everything was self-merged on green CI in isolated worktrees (siblings merged concurrently throughout, untouched). One thing worth noting for your records: **this is your real GPU box** (torch+CUDA/cv2/WSL/Gemini all live), not the GPU-less dev container the old CLAUDE.md note assumes ? I've flagged that in memory.

Want me to wire the full Gemini highlight run for GX010125 into the per-video layout and kick it off?

#### 370. user (2026-06-14T05:00:37.515Z)

yes please! And what were the %age improvements can you summarize and show

#### 371. assistant (2026-06-14T05:02:23.912Z)

Yes ? let me wire the full Gemini-confirmed run reusing the cached trajectory + Gate-A into the isolated `output/GX010125\_run/`, validate the cheap parts, then launch it in the background. First the runner:

#### 372. assistant (2026-06-14T05:02:45.651Z)

""""FULL Gemini-confirmed highlight run for GX010125\_1080p ? REUSES the cached WASB trajectory (no GPU re-detect), applies the promoted Gate-A (net-crossings >= 2), runs the Gemini handoff on the kept windows, and stitches the reel ? all from the 1080p, all into the per-video dir output/GX010125\_run/. Honors "reuse 1080p", isolated, no pollution.

This replicates trajectory\_hybrid's Phase-3/4 over the cached trajectory so we skip the ~tens-of-min WASB pass. Costs a little Gemini (\$); ~10-20 min wall.

""""

```
import concurrent.futures
import json
import logging
import os
import sys
import time
```

```
sys.path.insert(0, os.path.dirname(os.path.dirname(os.path.abspath(__file__))))
from dotenv import load_dotenv
```

```
load_dotenv()
logging.basicConfig(level=logging.INFO,
                    format="%(asctime)s [%(name)s] %(levelname)s: %(message)s",
                    datefmt="%H:%M:%S")
log = logging.getLogger("gx_full")
```

```
from backend.config import load_config
from backend.pipeline.detectors.tracknet_runner import TrackNetRunner
from backend.pipeline.detectors import rally_gate
from backend.providers import provider_registry
from backend.sports import sport_registry
from backend.utils.video import extract_video_chunk, get_video_duration
from backend.pipeline.stitcher.compiler import VideoCompiler
```

```
VIDEO = "output/GX010125_1080p.mp4"
TRAJ = "output/wasb/GX010125_1080p__89605e80ed01_wasb.csv"
OUT_DIR = "output/GX010125_run"
```

```
FPS, FRAME_W = 59.94, 1920.0
MIN_CROSSINGS = 2          # Gate-A (promoted)
```

```
def main():
    t0 = time.time()
    cfg = load_config()
    os.makedirs(OUT_DIR, exist_ok=True)
    handoff_dir = os.path.join(OUT_DIR, "handoff")
    os.makedirs(handoff_dir, exist_ok=True)
    idx = cfg.indexing
    vt, mg = idx.wasb.velocity_threshold, idx.dead_time_merge_gap
    mwd, pad = idx.min_action_window_duration, idx.ai_handoff_padding
    provider_name, model_name = cfg.ai_provider, cfg.ai_model
    log.info("config: provider=%s model=%s | windowing vt=%.3f merge_gap=%.1f min_dur=%.1f",
            pad, provider_name, model_name, vt, mg, mwd, pad)

    points = TrackNetRunner.parse_trajectory_csv(TRAJ)
    log.info("parsed %d cached trajectory points (%d visible)", len(points),
            sum(1 for p in points if p.visible))
    windows = TrackNetRunner.trajectory_to_action_windows(
        points, fps=FPS, frame_width=FRAME_W, chunk_start=0.0,
        velocity_thresh=vt, merge_gap=mg, min_window_duration=mwd, chunk_index=0)
    intervals = [[w["start"], w["end"]] for w in windows]
    kept = rally_gate.apply_gate(points, intervals, FPS, min_crossings=MIN_CROSSINGS)
    log.info("production windowing -> %d candidates; Gate-A(min_crossings=%d) -> %d kept",
            len(intervals), MIN_CROSSINGS, len(kept))

    api_key = os.environ.get(f"{provider_name.upper()}_API_KEY")
    if not api_key:
        raise SystemExit(f"{provider_name.upper()}_API_KEY not set in .env")
    provider = provider_registry.get(provider_name, api_key=api_key, model_name=model_name)
    sport = sport_registry.get(cfg.sport)
    duration = get_video_duration(VIDEO0)

    def handle(i, w):
        s = max(0.0, w[0] - pad)
        e = min(duration, w[1] + pad)
        clip = os.path.join(handoff_dir, f"hand_{i}.mp4")
        if not extract_video_chunk(VIDEO0, s, e - s, clip, reencode=True):
            return []
        prompt = sport.get_prompt(chunk_duration=e - s)
        segs, _ = provider.analyze_video(clip, prompt, chunk_duration=e - s, label=f"GX {i +
1}/{len(kept)}")
        out = []
        for sg in segs:
            try:
                out.append((float(sg["start_time"]) + s, float(sg["end_time"]) + s,
                    sg.get("shuttle_exchange_count")))
            except (ValueError, TypeError, KeyError):
                pass
        try:
            os.remove(clip)
        except OSError:
            pass
        log.info("window %d/%d -> %d rallies", i + 1, len(kept), len(out))
        return out

    confirmed = []
    with concurrent.futures.ThreadPoolExecutor(max_workers=idx.ai_max_workers) as ex:
        for fut in [ex.submit(handle, i, w) for i, w in enumerate(kept)]:
            confirmed.extend(fut.result())
    confirmed.sort()
    log.info("Gemini confirmed %d rally segments (%.0fs elapsed)", len(confirmed), time.time() -
```



```

t0)

with open(os.path.join(OUT_DIR, "GX010125_rallies.csv"), "w", encoding="utf-8") as f:
    f.write("start_sec,end_sec,shot_count\n")
    for s, e, sc in confirmed:
        f.write(f"{s:.3f},{e:.3f},{sc if sc is not None else ''}\n")

reel_path = None
rallies = [(s, e) for s, e, _ in confirmed]
if rallies:
    comp = VideoCompiler(OUT_DIR, begin_padding=cfg.stitching.begin_padding,
                        end_padding=cfg.stitching.end_padding,
watermark=cfg.stitching.watermark)
    reel_path = comp.compile_highlights(VIDEO, rallies, "GX010125_highlights.mp4")
    log.info("highlight reel: %s", os.path.abspath(reel_path))

    json.dump({"video": "GX010125_1080p", "source": "cached WASB trajectory (no GPU re-detect)",
              "candidates": len(intervals), "gate_a_kept": len(kept),
              "gemin_i_confirmed": len(confirmed), "reel": reel_path,
              "elapsed_sec": round(time.time() - t0)},
              open(os.path.join(OUT_DIR, "GX010125_full_summary.json"), "w"), indent=2)
try:
    os.rmdir(handoff_dir)
except OSError:
    pass
log.info("[DONE] %s rallies -> %s (%.0fs total)", len(rallies), OUT_DIR, time.time() - t0)

if __name__ == "__main__":
    main()

```

### 373. user (2026-06-14T05:02:45.679Z)

File created successfully at: C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/run\_gx010125\_full.py (file state is current in your context ? no need to Read it back)

### 374. assistant (2026-06-14T05:02:52.256Z)

Now validate the cheap parts (config, trajectory parse, production windowing, Gate-A, provider init, one clip extract) **\*\*before\*\*** spending any Gemini calls.

### 375. user (2026-06-14T05:03:16.093Z)

```

provider=gemini model=gemini-flash-latest
production windowing -> 18 candidates | Gate-A kept -> 16
key set: True
provider init OK: GeminiProvider
sport prompt OK len=3404
clip extract OK: True size= 129263462

```

### 376. assistant (2026-06-14T05:03:48.420Z)

All pieces validate ? **\*\*production windowing ? 18 candidates, Gate-A keeps 16\*\*** (note: production windowing already merges aggressively, so Gate-A drops fewer here than on the recall-oriented proposal set; an honest nuance). Provider + clip extract work. Launching the full run in the background (with keep-awake so Modern Standby doesn't kill it), logging to the isolated dir.

### 377. user (2026-06-14T05:03:49.503Z)

```

scratch/keep_aware.ps1
keep_aware present

```

### 378. user (2026-06-14T05:03:55.659Z)

keep-awake launched

### 379. user (2026-06-14T05:03:59.392Z)

Command running in background with ID: br3vux8hv. Output is being written to:  
C:\Users\avidu\AppData\Local\Temp\claudex\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\tasks\br3vux8hv.output. You will be notified when it completes. To check interim output, use Read on that file path.

### 380. assistant (2026-06-14T05:04:06.910Z)

The full run is launched in the background (reusing the cached trajectory + Gate-A ? Gemini ? reel, all into `output/GX010125\_run/`). Let me give it a moment, confirm it's progressing cleanly, then give you the percentage-improvement summary.

### 381. user (2026-06-14T05:04:38.945Z)

=== run log so far ===

```
10:34:17 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:17 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:18 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:18 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:19 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:19 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:20 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:20 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:21 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:21 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:25 [httpx] INFO: HTTP Request: POST https://generativelanguage.googleapis.com/upload/v1beta/files?upload_id=AJ5rDhG7D0mEja02siZaF2ItkNkxAd4quSczbeVaxxi0_bXrw1S0a1gwnmIjrkNR0mxobsLuJdr0LW
ORgh4RaEN5joNYoJ6nW0jBEWH4mVbGGHy&upload_protocol=resumable "HTTP/1.1 200 OK"
10:34:25 [backend.providers.gemini] INFO: [Gemini GX 1/16] Upload complete in 12.2s. URI:
https://generativelanguage.googleapis.com/v1beta/files/oqt942n9jfdm
10:34:25 [backend.providers.gemini] INFO: [Gemini GX 1/16] Processing on Google servers.
Elapsed: 0.0s. Retrying in 10s...
10:34:35 [httpx] INFO: HTTP Request: GET
https://generativelanguage.googleapis.com/v1beta/files/oqt942n9jfdm "HTTP/1.1 200 OK"
10:34:35 [backend.providers.gemini] INFO: [Gemini GX 1/16] Processing on Google servers.
Elapsed: 10.4s. Retrying in 10s...
```

### 382. assistant (2026-06-14T05:05:07.832Z)

The full run is healthy and progressing ? it parsed the cached trajectory, windowed (18 ? Gate-A kept 16), and is now in the **Gemini handoff** ( `[Gemini GX 1/16]`, clips uploading + processing). It'll land a real highlight reel in `output/GX010125\_run/` in ~10?15 min; I'll be notified when `br3vux8hv` finishes. (Background ID `br3vux8hv`, log at `output/GX010125\_run/full\_run.log`, keep-awake running.)

## Percentage improvements (measured this session)

All on the **n=6** golden corpus, offline LOV0 unless noted. Relative % = (new ? old) / old.

| Change                  | Metric                              | Before ? After            | Improvement                      |
|-------------------------|-------------------------------------|---------------------------|----------------------------------|
| Gate-A (promoted)       | mean F1 (tIoU 0.5)                  | 0.388 ? 0.462             | +19.1% (6/6 matches, p=0.031)    |
| Gate-A                  | F1@10 / F1@25                       | 0.645?0.748 / 0.603?0.698 | +16.0% / +15.8%                  |
| Gate-A                  | over-segmentation (seg-count ratio) | 1.68 ? 1.01               | -40% (excess fragmentation ~98%) |
| Learned fix (#130)      | mean F1                             | 0.307 ? 0.423             | +37.8%                           |
| Learned fix             | `testlarge` (was zeroed)            | 0.000 ? 0.300             | recovered from total failure     |
| Learned fix             | learned vs CONTROL                  | ?20.9% ? +9.0%            | flipped from losing to winning   |
| First single-video demo | F1 / precision (testlarge)          | 0.300?0.375 / 0.214?0.300 | +25.0% / +40.2%                  |

**Headline:** Gate-A is a **+19% F1** win across all 6 golden matches (and nearly eliminates over-segmentation); the threshold fix is a **+38% F1** rescue of the learned variant (and flips it from ?21% to +9% vs the heuristic baseline).

### Honest caveats:

- These are offline re-scores on the golden corpus's recall-oriented proposal candidate set. On the **production** (already-merged) windowing ? which the GX010125 run uses ? Gate-A is more surgical (drops 2/18 windows), because there's less fragmentation to catch. So Gate-A's % gain is candidate-set-dependent.
- **GX010125** itself has no golden labels, so its numbers are a **filter rate**, not an accuracy %: on the proposal windowing Gate-A dropped **18/45 (40%)** of candidate windows as likely held-shuttle/feed FPs; on production windowing, 2/18.
- The learned fix's gains are validated; per the documented history, the per-frame model wins by more at scale ? these window-level numbers are the floor.

I'll report the GX010125 reel (rally count, runtime, path) as soon as the background run completes.

### 383. user (2026-06-14T05:07:09.133Z)

```
<task-notification>
<task-id>br3vux8hv</task-id>
<tool-use-id>toolu_01LmUisvegVdyFe431SrsUG9</tool-use-id>
<output-file>C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-indexer\7d6bbb83-42c0-459b-9e23-e14e96d4cb2d\tasks\br3vux8hv.output</output-file>
<status>failed</status>
<summary>Background command "Run the full GX010125 Gemini highlight pipeline (background)" failed with exit code 127</summary>
</task-notification>
```

### 384. assistant (2026-06-14T05:07:21.458Z)

The run failed (exit 127) after progressing into the Gemini handoff ? let me read the full log to see where it broke.

### 385. user (2026-06-14T05:07:24.062Z)

```
=== tail of run log ===
10:34:01 [gx_full] INFO: config: provider=gemini model=gemini-flash-latest | windowing vt=0.020 merge_gap=2.0 min_dur=2.0 pad=2.0
10:34:01 [gx_full] INFO: parsed 46080 cached trajectory points (29879 visible)
```

```

10:34:01 [gx_full] INFO: production windowing -> 18 candidates; Gate-A(min_crossings=2) -> 16
kept
10:34:12 [backend.providers.gemini] INFO: [Gemini GX 1/16] Uploading
output/GX010125_run\handoff\hand_0.mp4...
10:34:25 [backend.providers.gemini] INFO: [Gemini GX 1/16] Upload complete in 12.2s. URI:
https://generativelanguage.googleapis.com/v1beta/files/oqt942n9jfdm
10:34:45 [backend.providers.gemini] INFO: [Gemini GX 1/16] Processing finished in 20.8s.
10:34:45 [backend.providers.gemini] INFO: [Gemini GX 1/16] Submitting multimodal request to
gemini-flash-latest...
10:34:45 [google_genai.models] INFO: AFC is enabled with max remote calls: 10.
=== summary/rallies written? ===
total 41
drwxr-xr-x 1 avidu 197609 0 Jun 14 10:34 .
drwxr-xr-x 1 avidu 197609 0 Jun 14 10:09 ..
-rw-r--r-- 1 avidu 197609 1082 Jun 14 09:54 GX010125_control_rallies.csv
-rw-r--r-- 1 avidu 197609 672 Jun 14 09:54 GX010125_gateA_rallies.csv
-rw-r--r-- 1 avidu 197609 306 Jun 14 09:54 GX010125_rescore_summary.json
-rw-r--r-- 1 avidu 197609 5561 Jun 14 10:34 full_run.log
drwxr-xr-x 1 avidu 197609 0 Jun 14 10:34 handoff
=== last 5 raw lines ===
10:34:35 [backend.providers.gemini] INFO: [Gemini GX 1/16] Processing on Google servers.
Elapsed: 10.4s. Retrying in 10s...
10:34:45 [httpx] INFO: HTTP Request: GET
https://generativelanguage.googleapis.com/v1beta/files/oqt942n9jfdm "HTTP/1.1 200 OK"
10:34:45 [backend.providers.gemini] INFO: [Gemini GX 1/16] Processing finished in 20.8s.
10:34:45 [backend.providers.gemini] INFO: [Gemini GX 1/16] Submitting multimodal request to
gemini-flash-latest...
10:34:45 [google_genai.models] INFO: AFC is enabled with max remote calls: 10.

```

### 386. assistant (2026-06-14T05:07:39.936Z)

The Python log ends cleanly mid-request (no traceback) then exit 127 ? that pattern means the process was \*killed\*, not a Python crash. Let me check the task output file for the shell-level cause.

### 387. user (2026-06-14T05:07:45.160Z)

```

=== task .output (shell-level) ===
=== leftover handoff clips (proof of where it stopped) ===
hand_0.mp4
=== is python still running? ===
python.exe                20888 Console                1  1,487,764 K

```

### 388. assistant (2026-06-14T05:08:23.383Z)

A python process (PID 20888, 1.5 GB) is still alive and `hand\_0.mp4` is mid-flight ? so either the run is orphaned-but-progressing, or stuck. Let me check whether the log is still growing.

### 389. user (2026-06-14T05:08:41.931Z)

```

now: 10:38:29
log last-modified: 10:34:45
log line count: 28
--- after 12s ---
log line count: 28
log last-modified: 10:34:45
10:34:45 [backend.providers.gemini] INFO: [Gemini GX 1/16] Processing finished in 20.8s.
10:34:45 [backend.providers.gemini] INFO: [Gemini GX 1/16] Submitting multimodal request to
gemini-flash-latest...
10:34:45 [google_genai.models] INFO: AFC is enabled with max remote calls: 10.

```

### 390. assistant (2026-06-14T05:09:56.035Z)

The run is \*\*hung\*\* ? log stale 4 min, stuck on the first Gemini `generateContent` (no client-

side timeout on that call), and the harness's background wrapper exited 127 leaving an orphaned hung python. I'll kill the orphan and make the runner robust (overall handoff timeout + force-exit past hung threads), then retry foreground so a single API hang can't block the whole run.

**391. user (2026-06-14T05:09:57.759Z)**

SUCCESS: The process with PID 20888 has been terminated.  
cleaned stale clips

**392. assistant (2026-06-14T05:10:07.113Z)**

Now make the runner hang-proof (overall handoff timeout + force-exit past stuck threads):

**393. user (2026-06-14T05:10:09.654Z)**

The file C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/run\_gx010125\_full.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**394. user (2026-06-14T05:10:17.155Z)**

The file C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/run\_gx010125\_full.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**395. user (2026-06-14T05:10:23.590Z)**

The file C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/run\_gx010125\_full.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**396. assistant (2026-06-14T05:10:42.053Z)**

Now retry foreground (bounded ~7 min so it can't hang the turn). First, a quick check whether the clip-extract helper can downscale ? large 1080p clips likely contributed to the generate stall.

**397. user (2026-06-14T05:10:43.982Z)**

```
20: def extract_video_chunk(video_path: str, start_time: float, duration: float,
def extract_video_chunk(video_path: str, start_time: float, duration: float,
                        output_path: str, reencode: bool = False,
                        scale_width: Optional[int] = None) -> bool:
    """Extracts a chunk from a video file.

    Args:
        reencode: If True, re-encodes with libx264 (slower but frame-accurate).
                  If False, uses stream-copy (fast but cuts at nearest keyframe).
                  Use True for short clips where keyframe drift matters.
        scale_width: If set (requires reencode), downscale to this width (height auto,
                    even). Useful for high-res sources headed to an API that analyses at
                    low res anyway ? shrinks uploads with no analysis loss.
    """
```

**398. assistant (2026-06-14T05:10:54.907Z)**

`scale\_width` is supported ? downscaling to 1280 (the configured `ai\_chunk\_scale\_width`) shrinks the ~129 MB clips ~10x and was likely the stall cause. Adding it, then retrying foreground.

**399. user (2026-06-14T05:10:56.988Z)**

The file C:/Users/avidu/Projects/badminton-highlight-indexer/scratch/run\_gx010125\_full.py has been updated successfully. (file state is current in your context ? no need to Read it back)

**400. user (2026-06-14T05:18:56.855Z)**

```

    },
    {
      "start_time": 9.5,
      "end_time": 21.0,
      "shuttle_exchange_count": 9,
      "shot_count_confidence": "High"
    },
    {
      "start_time": 27.0,
      "end_time": 35.0,
      "shuttle_exchange_count": 6,
      "shot_count_confidence": "High"
    },
    {
      "start_time": 41.0,
      "end_time": 45.0,
      "shuttle_exchange_count": 3,
      "shot_count_confidence": "High"
    }
  ]
}
10:48:26 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #9/18 [180.09s - 192.59s]
-> [179.592s - 193.092s] (duration: 13.50s)
10:48:28 [gx_full] INFO: window 5/16 -> 4 rallies
10:48:34 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #10/18 [195.05s - 196.85s]
-> [194.546s - 197.346s] (duration: 2.80s)
10:48:35 [backend.providers.gemini] INFO: [Gemini GX 6/16] Uploading
output/GX010125_run\handoff\hand_5.mp4...
10:48:36 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #11/18 [202.15s - 204.85s]
-> [201.646s - 205.346s] (duration: 3.70s)
10:48:38 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #12/18 [210.85s - 215.85s]
-> [210.346s - 216.346s] (duration: 6.00s)
10:48:40 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #13/18 [219.85s - 224.35s]
-> [219.346s - 224.846s] (duration: 5.50s)
10:48:42 [backend.providers.gemini] INFO: [Gemini GX 6/16] Upload complete in 6.3s. URI:
https://generativelanguage.googleapis.com/v1beta/files/milewf1e8gr6
10:48:43 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #14/18 [230.85s - 236.35s]
-> [230.346s - 236.846s] (duration: 6.50s)
10:48:45 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #15/18 [246.55s - 250.25s]
-> [246.047s - 250.747s] (duration: 4.70s)
10:48:48 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #16/18 [256.55s - 260.25s]
-> [256.047s - 260.747s] (duration: 4.70s)
10:48:50 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #17/18 [264.55s - 269.25s]
-> [264.047s - 269.747s] (duration: 5.70s)
10:48:52 [backend.pipeline.stitcher.compiler] INFO: Trimming Segment #18/18 [275.75s - 283.25s]
-> [275.247s - 283.747s] (duration: 8.50s)
10:48:52 [backend.providers.gemini] INFO: [Gemini GX 6/16] Processing finished in 10.5s.
10:48:52 [backend.providers.gemini] INFO: [Gemini GX 6/16] Submitting multimodal request to
gemini-flash-latest...
10:48:56 [backend.pipeline.stitcher.compiler] INFO: Successfully trimmed 18/18 clips. Preparing
to stitch...
10:48:56 [backend.pipeline.stitcher.compiler] INFO: Stitching clips together...
10:48:56 [backend.pipeline.stitcher.compiler] INFO: Dynamic highlight reel saved:
output/GX010125_run\GX010125_highlights.mp4 (156.6 MB)
10:48:56 [gx_full] INFO: highlight reel: C:\Users\avidu\Projects\badminton-highlight-
indexer\output\GX010125_run\GX010125_highlights.mp4
10:48:56 [gx_full] INFO: [DONE] 18 rallies -> output/GX010125_run (472s total)

```

#### 401. assistant (2026-06-14T05:19:23.619Z)

The run **\*\*completed\*\*** this time (downscaling to 1280 + the timeout safety net fixed the stall) ? **\*\*18 rallies stitched into `GX010125\_highlights.mp4` (156.6 MB)\*\***. Let me verify the isolated outputs and how many windows actually confirmed (the overall handoff timeout fired at 420s, so it may be a partial set).

#### 402. user (2026-06-14T05:19:28.946Z)

```

=== full run summary ===
{
  "video": "GX010125_1080p",
  "source": "cached WASB trajectory (no GPU re-detect)",
  "candidates": 18,
  "gate_a_kept": 16,
  "gemini_confirmed": 18,
  "reel": "output/GX010125_run\\GX010125_highlights.mp4",
  "elapsed_sec": 472
}
=== confirmed rallies (count + first few) ===
19
start_sec,end_sec,shot_count
6.800,15.500,6
66.992,68.592,2
71.592,80.592,10
82.092,84.592,3
88.592,90.592,2
=== isolated dir (per-video layout) ===
total 160443
drwxr-xr-x 1 avidu 197609      0 Jun 14 10:48 .
drwxr-xr-x 1 avidu 197609      0 Jun 14 10:09 ..
-rw-r--r-- 1 avidu 197609    1082 Jun 14 09:54 GX010125_control_rallies.csv
-rw-r--r-- 1 avidu 197609     243 Jun 14 10:48 GX010125_full_summary.json
-rw-r--r-- 1 avidu 197609     672 Jun 14 09:54 GX010125_gateA_rallies.csv
-rw-r--r-- 1 avidu 197609 164246446 Jun 14 10:48 GX010125_highlights.mp4
-rw-r--r-- 1 avidu 197609     363 Jun 14 10:48 GX010125_rallies.csv
-rw-r--r-- 1 avidu 197609     306 Jun 14 09:54 GX010125_rescore_summary.json
-rw-r--r-- 1 avidu 197609    5561 Jun 14 10:34 full_run.log
=== past GX010125 work untouched? ===
output/GX010125_highlights.mp4 1207870728
=== stop keep-awake ===
keep-awake stopped (best-effort)

```

#### 403. assistant (2026-06-14T05:19:49.441Z)

The reel is built and **past work is intact** (the old 1.2 GB 4K reel untouched). But the handoff hit its 420s cap (472s total > 420s), so I need to verify whether all 16 Gate-A windows confirmed or only a subset before the timeout ? to report honestly.

#### 404. user (2026-06-14T05:19:53.674Z)

```

=== per-window completions logged ===
=== count of windows that returned ===
0
=== overall-timeout fired? ===
=== rally total duration (min) ===
1.6 min across 18 rallies

```

#### 405. assistant (2026-06-14T05:20:51.394Z)

The full Gemini-confirmed run **succeeded** ? here's the honest result.

#### GX010125 highlight reel ? (`output/GX010125\_run/`)

- **GX010125\_highlights.mp4** ? 18 rallies, ~1.6 min, 156 MB, stitched from the 1080p.
- Reused the **cached 1080p WASB trajectory** (no GPU re-detect, no 4K); applied the promoted **Gate-A** (18 candidates ? 16 kept); Gemini confirmed the rallies; everything landed in the **isolated per-video dir**.
- **Past work intact** ? the old 1.2 GB 4K reel (`output/GX010125\_highlights.mp4`) was not touched.

**Honest caveat:** the first attempt **hung** on a Gemini `generateContent` call (no client-side timeout; the 1080p clips were ~129 MB each). I fixed it two ways ? **downscaling clips to 1280** (uploads dropped 12s?6s) and a **420s overall handoff cap + force-exit** so a stuck call can't

block the run. That cap \*did\* fire (run was 472s), so the reel is the rallies that confirmed within ~7 min ? likely a \*\*subset\*\* of the 16 windows (the older 4K run got 52 rallies over ~47 min). For exhaustive coverage, bump `HANDOFF\_TIMEOUT` in `scratch/run\_gx010125\_full.py` (generate runs ~30?60s/window). Cleanup done: orphaned process killed, keep-awake stopped.

## Everything this turn ? all merged / delivered

```
| Item | Result |
|---|---|
| **Fix the logistic** ([#130](https://github.com/avidullu/badminton-highlight-indexer/pull/130)) | tune_threshold + calibrate ? mean F1 **0.307?0.423 (+38%)**, testlarge un-zeroed, beats CONTROL |
| **Promote Gate-A** ([#131](https://github.com/avidullu/badminton-highlight-indexer/pull/131)) | doc trail + recommended config; live flip owner-gated |
| **Per-video layout plan** ([#131](https://github.com/avidullu/badminton-highlight-indexer/pull/131)) | `docs/PER_VIDEO_OUTPUT_LAYOUT.md` (fixes the hardcoded `output/` paths) |
| **Run GX010125** | real reel, isolated, reused 1080p (above) |
```

## % improvements (recap)

- \*\*Gate-A:\*\* +19.1% mean F1 (6/6 golden matches,  $p=0.031$ ); over-segmentation ?40%.
- \*\*Logistic fix:\*\* +37.8% mean F1; flipped from ?21% to +9% vs the heuristic.
- \*\*GX010125 (no labels):\*\* Gate-A filtered 40% of proposal candidates / 2 of 18 production windows.

One forward note worth your attention: the Gemini handoff has \*\*no client-side timeout on the generate call\*\* (it hangs rather than retries) ? that's a real production fragility I'd suggest hardening in `backend/providers/gemini.py` as a small follow-up. Want me to spin that up?

### 406. user (2026-06-14T05:30:31.087Z)

did you generate a new stitched highlight for gx010125? And yes, please do start that work and make sure its a policy driven system and the default is WAIT\_AND\_RESUME with ABORT, POLL\_EVERY\_N\_SEC etc. like relevant policies can be applied? I believe each job can specify its own and can self schedule? So a master node etc. would not be necessary?

### 407. assistant (2026-06-14T05:31:54.367Z)

Yes ? I did generate a real new reel for GX010125:  
\*\*`output/GX010125\_run/GX010125\_highlights.mp4`\*\* (18 Gemini-confirmed rallies, ~1.6 min, 156 MB), built from the 1080p with Gate-A applied, isolated, with the old 4K reel untouched. (Caveat from my last message: it's likely a \*subset\* because the 7-min safety cap fired ? exactly the fragility you're now asking to fix properly.)

And yes ? let me start that work the right way: a \*\*policy-driven\*\* design (not an ad-hoc timeout). Your architecture instincts are sound; let me validate them against what's already in the repo before spec'ing. Let me read the existing async-job model and the exact Gemini hang point.

### 408. user (2026-06-14T05:31:59.112Z)

=== gemini provider: the generate/poll/retry surface (where the hang is) ===

```
29:     def analyze_video(self, video_path: str, prompt: str,
56:         for attempt in range(3):
58:             video_file = self.client.files.upload(file=video_path)
65:             time.sleep(2 * (attempt + 1))
75:             while video_file.state.name == "PROCESSING":
84:                 time.sleep(10)
85:                 video_file = self.client.files.get(name=video_file.name)
108:         # 6 attempts with exponential backoff + jitter (~3s..96s, worst case ~3min total):
112:         for attempt in range(gen_attempts):
114:             response = self.client.models.generate_content(
129:                 time.sleep(3 * (2 ** attempt) + random.uniform(0, 2))
```

=== async job model: executor + job lifecycle in main.py ===



```

8:from concurrent.futures import ThreadPoolExecutor
71:_job_executor: Optional[ThreadPoolExecutor] = None
73:def _get_executor() -> ThreadPoolExecutor:
77:    _job_executor = ThreadPoolExecutor(max_workers=1, thread_name_prefix="jobworker")
238:def _run_job(job_id: str, req: ProcessRequest) -> None:
241:    Job `status` is the EXECUTION state (queued|running|done|failed): 'done' means the
259:    """Submit a processing job (202 + job_id). Poll GET /api/jobs/{job_id} for state.
286:    content={"job_id": job_id, "status": "queued", "poll": f"/api/jobs/{job_id}"},
290:@app.get("/api/jobs/{job_id}")
687:    # Crash recovery (#16): jobs left queued/running by a previous process are dead ?
689:    swept = db_instance.fail_stale_jobs()

=== DEFERRED_HARDENING_PLAN async-job (#16) section anchors ===
13:#13 Train/eval guardrails (correctness of the data flywheel) ? #16 Async jobs
90:## 3. Async job model (debt #16)
95:> polls. See `scratch/PR1_HANDOFF.md` for the review bundle.
110:- Update the static UI to poll `GET /api/jobs/{id}` instead of blocking.
121:- Submit a job ? poll ? DONE with the same segments the sync path produced.

```

#### 409. assistant (2026-06-14T05:32:57.055Z)

Now I have the picture. The two unbounded waits are in `backend/providers/gemini.py`: the  
**\*\*`while state == "PROCESSING": sleep(10)` poll loop\*\*** (no max-wait) and the  
**\*\*`generate\_content` call\*\*** (no per-call timeout ? the retry loop only catches *\*failures\**, not *\*hangs\**). And the async-job substrate (single-worker executor + `jobs` table + crash-recovery sweep, DEFERRED\_HARDENING #16) already has **\*\*no master node\*\*** ? your instinct is correct. Let me spec this as a policy framework. Worktree + design doc:

#### 410. user (2026-06-14T05:33:01.452Z)

c8bdfdd docs: promote Gate-A (measured win) + per-video output layout plan (#131)  
 HEAD is now at c8bdfdd docs: promote Gate-A (measured win) + per-video output layout plan (#131)

#### 411. assistant (2026-06-14T05:33:46.063Z)

## Execution policy ? per-job, self-scheduling resilience for long/flaky operations

**\*\*Status:\*\*** **\*\*DESIGN ? owner-gated, NOT started.\*\*** Motivated by a real GX010125 hang: the Gemini  
 `generateContent` call stalled with **\*\*no client-side timeout\*\*** and blocked the whole run. The  
 fix is  
 not a one-off timeout ? it is a **\*\*policy\*\*** every long operation carries, so waiting / polling /  
 aborting /  
 resuming is declared per job, not hard-coded. Docs-only.

### The problem (concrete)

`backend/providers/gemini.py::analyze\_video` has **\*\*two unbounded waits\*\***:  
 1. the file-processing poll ? `while video\_file.state.name == "PROCESSING": time.sleep(10)` (no  
 deadline);  
 2. the `generate\_content` call ? the retry loop (6 attempts, backoff) catches **\*\*failures\*\***, not  
**\*\*hangs\*\***;  
 a request that never returns blocks forever.  
 The same shape recurs anywhere we wait on an external/long thing (WSL/WASB GPU pass, ffmpeg, a  
 future  
 remote compute backend). Today each site invents its own ad-hoc loop (or none). We want one  
 declarative  
 contract.

### Your architecture ? validated

- **\*\*"Each job can specify its own [policy]."** ? Yes ? an `ExecutionPolicy` is attached **\*\*per  
 job\*\***  
 (and per long sub-operation), defaulting sensibly, overridable by the caller. It is just data  
 on the job.  
 - **\*\*"Jobs can self-schedule."** ? Yes ? a job that must wait persists its `next\_poll\_at` +

progress and

**\*\*re-enqueues itself\*\***; the existing single worker re-claims it when due. No periodic central tick.

- **\*\*"A master node is not necessary."\*\*** ? Correct. The repo is already master-less: a

``ThreadPoolExecutor``

worker (``main.py``, ``max_workers=1``, single self-hosted box) drains the ``jobs`` table, and a

**\*\*crash-recovery**

sweep (``fail_stale_jobs()``, DEFERRED\_HARDENING #16) reaps jobs a dead process left behind.

Self-scheduling

= cooperative re-enqueue; **\*\*any\*\*** worker can claim a due job, so it generalizes to N workers (PLATFORM\_ARCHITECTURE pluggable compute) **\*\*without\*\*** a dedicated scheduler ? the DB **\*is\*** the coordination.

## The contract

```
```python
```

```
class WaitMode(str, Enum):
```

```
    WAIT_AND_RESUME = "wait_and_resume" # DEFAULT ? bounded wait; on deadline, persist +
```

```
reschedule
```

```
    ABORT = "abort" # on deadline/hang, stop this op (fail or partial)
```

```
    BLOCK = "block" # legacy in-process bounded blocking wait (no
```

```
reschedule)
```

```
@dataclass(frozen=True)
```

```
class ExecutionPolicy:
```

```
    mode: WaitMode = WaitMode.WAIT_AND_RESUME
```

```
    poll_every_n_sec: float = 10.0 # cadence for polling an in-progress external op
```

```
    op_timeout_sec: float = 120.0 # HARD per-attempt deadline ? the hang-killer (was
```

```
missing)
```

```
    max_wait_sec: float = 1800.0 # total wall budget across resumes before giving up
```

```
    max_retries: int = 5 # failures (not hangs); exponential backoff + jitter
```

```
    backoff_base_sec: float = 3.0
```

```
    on_timeout: Literal["reschedule", "partial", "fail"] = "reschedule"
```

```
    resumable: bool = True # persist partial progress so a resume continues, not
```

```
restarts
```

```
    # JSON-serializable ? rides on the job row; a variant/job overrides any field.
```

```
...
```

- **\*\*`op\_timeout\_sec`\*\*** is the missing piece that would have killed the GX010125 hang: every external call

(upload, ``files.get``, ``generate_content``) runs under a hard per-attempt deadline; on expiry the policy's

``on_timeout`` decides ? ``reschedule`` (WAIT\_AND\_RESUME), ``partial`` (proceed with what's done), or ``fail``.

- **\*\*Hang vs failure are distinct.\*\*** Retries (``max_retries`` + backoff) handle a call that

**\*errors\***; the

**\*timeout\*** handles a call that **\*hangs\***. Both are policy, not provider-baked.

## Self-scheduling mechanism (no master)

A long op under ``WAIT_AND_RESUME``:

1. runs the external call under ``op_timeout_sec``;

2. on success ? record progress, continue;

3. on timeout/in-progress ? **\*\*persist progress + `next\_poll\_at = now + poll\_every\_n\_sec`**, set job

``status=waiting``, release the worker (don't hold a thread sleeping for minutes);

4. the worker loop (and the startup sweep) claims jobs whose ``next_poll_at <= now`` and resumes them;

5. give up when cumulative wait > ``max_wait_sec`` ? ``on_timeout`` terminal action.

This needs **\*\*only\*\*** two ``jobs``-table columns (``next_poll_at``, ``progress`` JSON ? ``progress`` already exists)

and a "claim due jobs" query. No timer service, no leader, no master. On one box it is one worker draining a

priority-by-due-time queue; on many boxes it is the same query with a row-claim (``UPDATE ? WHERE status=`

``waiting' AND next_poll_at<=now``), so a job migrates to whoever is free.

## Resumability (why WAIT\_AND\_RESUME is the default)

The op persists **what it already finished** so a resume does **not** redo it. For the AI handoff that means

**per-window confirmation state**: confirmed windows are written as they complete; a resume processes only the unconfirmed ones. The GX010125 run is the worked example ? instead of the 7-min cap abandoning ~10 windows and losing them, WAIT\_AND\_RESUME would persist the confirmed windows and reschedule the rest, so the reel becomes exhaustive across resumes with **no wasted Gemini spend**. ``ABORT`` (or ``on_timeout=partial``) is the opt-out for "good enough now" (what the one-off script does today).

## Phasing (owner-gated; ONE PR per slice, off `master`, default = current behaviour)

- **P0** ? this doc.
- **P1** ? ``ExecutionPolicy`` + bound the two Gemini waits (the actual hang fix). Add the dataclass; thread a policy into ``GeminiProvider.analyze_video``; run upload / ``files.get`` poll / ``generate_content`` under ``op_timeout_sec`` (a thread-future wrapper or the genai client's request timeout) with the poll deadline. Default policy reproduces today's behaviour **except** hangs now terminate per ``on_timeout`` instead of blocking forever. Pure-logic, unit-testable (a fake provider that hangs ? policy aborts/reschedules).
- **P2** ? per-job policy + self-scheduling on the ``jobs`` table (``next_poll_at``, claim-due query, ``status=waiting``); the executor honours WAIT\_AND\_RESUME by releasing + rescheduling instead of blocking.
- **P3** ? resumable AI handoff (persist per-window confirmation; resume processes only unconfirmed) ? makes the GX010125-style run exhaustive across resumes.
- **P4** ? apply the same policy to the WSL/WASB GPU pass + ffmpeg (the other long ops).

## Trade-offs / notes

- **Default WAIT\_AND\_RESUME, but bounded** ? ``max_wait_sec`` guarantees forward progress / eventual give-up; it never waits truly forever (the current bug).
- **Idempotent resumes** ? a resumed op must not double-spend (e.g. re-upload a clip already confirmed); P3's per-window state is what makes that safe.
- **Cost control** ? ``op_timeout_sec`` + ``partial`` is the cheap path; ``WAIT_AND_RESUME`` trades wall-clock for completeness without re-spending. Both are one field on the job.
- **No new infra** ? rides the existing ``jobs`` table + single-worker executor + crash sweep; the DB is the coordinator, so it scales to multi-worker (PLATFORM\_ARCHITECTURE) with no master added.

**Links:** the async-job substrate is ``docs/DEFERRED_HARDENING_PLAN.md`` §3 (#16); multi-worker future is ``docs/PLATFORM_ARCHITECTURE.md``; the motivating hang is the GX010125 run (``scratch/run_gx010125_full.py``, where a 7-min cap had to be hand-rolled).

### 412. user (2026-06-14T05:33:46.220Z)

File created successfully at: C:/Users/avidu/Projects/bhi-policy-wt/docs/EXECUTION\_POLICY.md (file state is current in your context ? no need to Read it back)

### 413. assistant (2026-06-14T05:33:59.229Z)

Now index it in ``docs/README.md``, then commit + self-merge on green.

### 414. user (2026-06-14T05:34:03.925Z)

<tool\_use\_error>File has not been read yet. Read it first before writing to it.</tool\_use\_error>

#### 415. user (2026-06-14T05:34:10.069Z)

11 - [DEFERRED\_HARDENING\_PLAN.md](DEFERRED\_HARDENING\_PLAN.md) ? Design (pending approval) for the larger deferred items: packaging, train/eval guardrails, and the async job model (#16, the first infra slice toward the platform).  
12 - [I18N\_PLAN.md](I18N\_PLAN.md) ? Design (owner-gated) for internationalizing the product \*\*by design\*\*: Kannada-first, Hindi next, extensible to other Indian languages. Shared framework-agnostic locale foundation across UI chrome, coach-facing backend text, and the future NL-query / coach-mode surfaces.

#### 416. user (2026-06-14T05:34:18.140Z)

The file C:/Users/avidu/Projects/bhi-policy-wt/docs/README.md has been updated successfully.  
(file state is current in your context ? no need to Read it back)

#### 417. user (2026-06-14T05:34:39.958Z)

warning: in the working copy of 'docs/EXECUTION\_POLICY.md', LF will be replaced by CRLF the next time Git touches it  
\* [new branch] docs/execution-policy -> docs/execution-policy  
<https://github.com/avidullu/badminton-highlight-indexer/pull/132>

#### 418. assistant (2026-06-14T05:34:45.664Z)

PR [#132](<https://github.com/avidullu/badminton-highlight-indexer/pull/132>). Watching CI, then self-merge + cleanup.

#### 419. user (2026-06-14T05:36:31.850Z)

Syntax sweep + import smoke (no GPU stack) pass 18s  
<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489588915/job/81252047726>  
Full test suite (offline, mocked) pass 1m33s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489588915/job/81252047704>  
Mypy (lenient baseline) pass 31s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489588915/job/81252047722>  
Ruff lint pass 9s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489588915/job/81252047713>  
Syntax sweep + import smoke (no GPU stack) pass 18s  
<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489588915/job/81252047726>

#### 420. user (2026-06-14T05:36:39.903Z)

MERGED  
worktree removed  
Deleted branch docs/execution-policy (was 153fa88).

#### 421. assistant (2026-06-14T05:37:02.146Z)

Done. Two answers:

##### 1. The GX010125 reel ? yes, generated ?

`output/GX010125\_run/GX010125\_highlights.mp4` ? \*\*18 Gemini-confirmed rallies, ~1.6 min, 156 MB\*\*, built from the 1080p with Gate-A, isolated, old 4K reel untouched. (Subset, because the hand-rolled 7-min cap fired ? exactly what the policy work below fixes.)

##### 2. Started the resilience work ? policy-driven, your architecture validated ? [#132](<https://github.com/avidullu>)

New \*\*\*`docs/EXECUTION\_POLICY.md`\*\*\*. I checked your hypotheses against the actual code and they hold:

- \*\*\*"Each job specifies its own policy"\*\*\* ? ? an `ExecutionPolicy` is just data on the job (JSON-serializable), defaulted and overridable.

- **Jobs self-schedule** ? ? a waiting job persists ``next_poll_at`` + progress and **re-enqueues itself**; the worker re-claims it when due. No periodic central tick.

- **No master node needed** ? **confirmed** ? the repo is **already** master-less: a single-worker ``ThreadPoolExecutor`` drains the ``jobs`` table + a crash-recovery sweep (`DEFERRED_HARDENING #16`). The **DB** is the coordinator, so it scales to N workers (a row-claim query) with no scheduler added.

**The contract** (your terms, plus the missing piece): ``WAIT_AND_RESUME`` (default) | ``ABORT`` | ``BLOCK``, with ``poll_every_n_sec``, ``op_timeout_sec`` (the hard per-attempt deadline that was missing and caused the GX010125 hang), ``max_wait_sec``, ``max_retries``, ``on_timeout`` (`reschedule|partial|fail`), ``resumable``. Key insight baked in: **a hang ? a failure** ? retries handle errors, the timeout handles stalls; today's code only had the former.

**The real fix** is precisely located: two unbounded waits in ``backend/providers/gemini.py`` ? the ``while state=="PROCESSING": sleep(10)`` poll loop and the ``generate_content`` call. **P1** wraps both in the policy (default reproduces today's behaviour **except** hangs now terminate per ``on_timeout`` instead of blocking forever). **P3** (resumable handoff: persist per-window confirmation) is what would make the GX010125 reel **exhaustive** across resumes with zero wasted Gemini spend, instead of the timeout abandoning windows.

Phased: **P1** = ``ExecutionPolicy`` + bound the two Gemini waits (the actual hang fix) . **P2** = per-job policy + self-scheduling on the ``jobs`` table . **P3** = resumable handoff . **P4** = WSL/WASB + ffmpeg.

Want me to implement **P1** now ? the ``ExecutionPolicy`` dataclass + bounding the two Gemini waits (the concrete hang fix), additive/default-safe, with unit tests (a fake provider that hangs ? policy aborts/reschedules)? That's the first code slice and it's self-contained.

#### 422. user (2026-06-14T05:38:46.436Z)

please go ahead since this is very important, ensure to document the decisions and complete the entire execution with self merge well distinct PRs. Testing should ensure very high coverage because this can lead to significant delays if not done well. Also do make sure to plumb logging so that the logs can tell that the jobs are polled, without polluting it!

#### 423. assistant (2026-06-14T05:40:23.143Z)

This is important and you want it done thoroughly ? distinct PRs per phase, self-merged, high coverage, clean poll-logging. I'll execute the phases sequentially. Let me start with **P1** (the ``ExecutionPolicy`` module + the actual Gemini hang fix). First I need the exact gemini surface to bound it correctly.

#### 424. user (2026-06-14T05:40:24.097Z)

```
1      import os
2      import time
3      import random
4      import json
5      import logging
6      from typing import Tuple
7      from google import genai
8      from google.genai import types
9      from backend.providers.base import AIVideoProvider
10
11     logger = logging.getLogger(__name__)
12
13     # The Gemini File API caps uploads at 2GB; stay well under it. Callers must chunk so
14     # each uploaded clip is below this. Oversize uploads are refused (not silently
15     # attempted).
16     MAX_UPLOAD_MB = 500
17
18     class GeminiProvider(AIVideoProvider):
19         def __init__(self, api_key: str, model_name: str):
20             super().__init__(api_key, model_name)
```

```

20         # Initialize client lazily to avoid connection/auth issues during class
instantiation
21         self._client = None
22
23     @property
24     def client(self):
25         if self._client is None:
26             self._client = genai.Client(api_key=self.api_key)
27         return self._client
28
29     def analyze_video(self, video_path: str, prompt: str,
30                      chunk_duration: float = 0.0,
31                      label: str = "") -> Tuple[list, dict]:
32         log_prefix = f"Gemini {label}" if label else "Gemini"
33
34         metrics = {
35             "upload_time": 0.0,
36             "process_time": 0.0,
37             "gen_time": 0.0
38         }
39
40         # 0. Refuse oversize uploads (Gemini File API limit is 2GB; we cap lower for
safety/speed)
41         try:
42             size_mb = os.path.getsize(video_path) / (1024 * 1024)
43         except OSError:
44             size_mb = 0.0
45         if size_mb > MAX_UPLOAD_MB:
46             logger.error(f"[{log_prefix}] {video_path} is {size_mb:.0f}MB >
{MAX_UPLOAD_MB}MB cap ? "
47                        f"SKIPPED WITHOUT ANALYSIS (no segments for this time range). "
48                        f"Re-encode/downscale chunks (indexing.ai_chunk_scale_width) or
chunk smaller.")
49             metrics["oversize_skipped"] = True # callers aggregate this into the run
report
50             return [], metrics
51
52         # 1. Upload (retry transient network failures, e.g. WinError 10053 / reset
connections)
53         logger.info(f"[{log_prefix}] Uploading {video_path}...")
54         t_upload_start = time.time()
55         video_file = None
56         for attempt in range(3):
57             try:
58                 video_file = self.client.files.upload(file=video_path)
59                 metrics["upload_time"] = time.time() - t_upload_start
60                 logger.info(f"[{log_prefix}] Upload complete in
{metrics['upload_time']:.1f}s. URI: {video_file.uri}")
61                 break
62             except Exception as e:
63                 logger.warning(f"[{log_prefix}] Upload attempt {attempt + 1}/3 failed:
{e}")
64                 if attempt < 2:
65                     time.sleep(2 * (attempt + 1))
66         if video_file is None:
67             logger.error(f"[{log_prefix}] Upload failed after 3 attempts.")
68             return [], metrics
69
70         # 2. Wait for processing on Google servers (bounded ? a stuck file must not hang
71         # the worker thread / request forever).
72         t_process_start = time.time()
73         process_deadline_s = 1200 # 20 min
74         try:
75             while video_file.state.name == "PROCESSING":
76                 if time.time() - t_process_start > process_deadline_s:

```

```

77             logger.error(f"[{log_prefix}] Processing exceeded
{process_deadline_s}s ? giving up.")
78             try:
79                 self.client.files.delete(name=video_file.name)
80             except Exception:
81                 pass
82             return [], metrics
83             logger.info(f"[{log_prefix}] Processing on Google servers. Elapsed:
{time.time()-t_process_start:.1f}s. Retrying in 10s...")
84             time.sleep(10)
85             video_file = self.client.files.get(name=video_file.name)
86             if video_file.state.name == "FAILED":
87                 logger.error(f"[{log_prefix}] Video processing failed:
{video_file.error.message}")
88             try:
89                 self.client.files.delete(name=video_file.name)
90             except Exception:
91                 pass
92             return [], metrics
93         except Exception as e:
94             logger.error(f"[{log_prefix}] Status check failed: {e}")
95             try:
96                 self.client.files.delete(name=video_file.name)
97             except Exception:
98                 pass
99             return [], metrics
100
101         metrics["process_time"] = time.time() - t_process_start
102         logger.info(f"[{log_prefix}] Processing finished in
{metrics['process_time']:.1f}s.")
103
104         # 3. Call Gemini (retry transient 503/UNAVAILABLE model-overload + network
blips)
105         logger.info(f"[{log_prefix}] Submitting multimodal request to
{self.model_name}...")
106         t_gen_start = time.time()
107         raw_text = None
108         # 6 attempts with exponential backoff + jitter (~3s..96s, worst case ~3min
total):
109         # capacity 503s on the flagship flash model come in waves that outlast the old
110         # 3x(3..9s) schedule ? observed live: 5/7 chunks of a run lost to three-strike
503s.
111         gen_attempts = 6
112         for attempt in range(gen_attempts):
113             try:
114                 response = self.client.models.generate_content(
115                     model=self.model_name,
116                     contents=[video_file, prompt],
117                     config=types.GenerateContentConfig(
118                         response_mime_type="application/json",
119                         temperature=0.0
120                     )
121                 )
122                 raw_text = response.text
123                 metrics["gen_time"] = time.time() - t_gen_start
124                 logger.info(f"[{log_prefix}] Response received in
{metrics['gen_time']:.1f}s:\n{raw_text}")
125                 break
126             except Exception as e:
127                 logger.warning(f"[{log_prefix}] generate attempt {attempt +
1}/{gen_attempts} failed: {e}")
128                 if attempt < gen_attempts - 1:
129                     time.sleep(3 * (2 ** attempt) + random.uniform(0, 2))
130             try:
131                 self.client.files.delete(name=video_file.name)

```

```

132         except Exception:
133             pass
134
135         if not raw_text:
136             return [], metrics
137
138         # 4. Parse JSON
139         try:
140             segments = json.loads(raw_text)

```

#### 425. assistant (2026-06-14T05:42:04.127Z)

The real unbounded wait is the `generate_content` call (the retry loop only catches \*exceptions\*, not a hang ? the processing-poll is already bounded at 1200s but spams INFO every 10s). P1 = an `ExecutionPolicy` module (`run_bounded` for the hang, `poll_until` for clean bounded polling) + refactor both gemini waits onto it, with fully injectable clocks so tests need zero real waiting. Let me build it.

#### 426. user (2026-06-14T05:42:09.784Z)

2673eee docs: execution-policy framework ? per-job, self-scheduling resilience (design) (#132)  
 HEAD is now at 2673eee docs: execution-policy framework ? per-job, self-scheduling resilience (design) (#132)  
 === existing gemini tests ===  
 test\_annotation\_providers.py  
 test\_automated\_provider.py  
 test\_gemini.py  
 test\_gemini\_refine.py  
 test\_provider\_gemini.py  
 test\_providers.py

#### 427. user (2026-06-14T05:42:43.437Z)

File created successfully at:  
 C:/Users/avidu/Projects/bhi-p1-wt/backend/infrastructure/execution\_policy.py (file state is current in your context ? no need to Read it back)

#### 428. assistant (2026-06-14T05:42:15.304Z)

There are gemini tests to keep green. Let me write the policy module first.

#### 429. assistant (2026-06-14T05:42:43.407Z)

"""Per-job execution policy: bounded, observable waits for long/flaky ops (EXECUTION\_POLICY.md P1).

Turns an unbounded blocking call (the GX010125 Gemini hang) into a bounded one:  
 - `run_bounded` ? a HARD per-attempt timeout (`op_timeout_sec`, the hang-killer) + retry-with-backoff  
   for FAILURES (an exception) \*and\* HANGS (no return in time). A hung attempt's worker thread is a daemon,  
   so it is abandoned, not joined ? it never blocks the process.  
 - `poll_until` ? a bounded poll loop for an in-progress external op (`poll_every_n_sec` cadence,  
   `max_wait_sec` deadline).

**Testability:** `sleep`/`monotonic` are injected, so the whole module is unit-tested with zero real waiting (a fake clock advances instantly).

**Logging discipline** (so polling is VISIBLE but never NOISY): every poll/retry logs at `DEBUG`; state transitions (start, resolved, hung, gave-up) log at `INFO/WARNING`. INFO stays clean; enable DEBUG on the



```

``execution.policy`` logger to watch the poll trail.
"""
from __future__ import annotations

import logging
import threading
import time as _time
from dataclasses import asdict, dataclass
from enum import Enum
from typing import Any, Callable, Optional, TypeVar

T = TypeVar("T")

#: Dedicated logger so callers can dial polling visibility independently of the app's root
level.
LOG = logging.getLogger("execution.policy")

class WaitMode(str, Enum):
    """How a job treats a long wait. ``str``-valued so it serialises straight to JSON on the job
row."""
    WAIT_AND_RESUME = "wait_and_resume" # default: bounded wait; over budget ->
reschedule/partial (P2/P3)
    ABORT = "abort" # one attempt; on hang/fail, stop immediately
    BLOCK = "block" # bounded in-process blocking wait (no reschedule)

class PolicyTimeout(TimeoutError):
    """An op hung past ``op_timeout_sec`` or never became ready within ``max_wait_sec``."""

@dataclass(frozen=True)
class ExecutionPolicy:
    """Declarative resilience knobs, attachable per job/op (JSON-serialisable; see
EXECUTION_POLICY.md)."""
    mode: WaitMode = WaitMode.WAIT_AND_RESUME
    poll_every_n_sec: float = 10.0
    op_timeout_sec: float = 120.0 # HARD per-attempt deadline ? the hang-killer that was
missing
    max_wait_sec: float = 1800.0 # total wall budget for a poll loop
    max_retries: int = 5 # retries (=> max_retries+1 attempts) for
failures/hangs
    backoff_base_sec: float = 3.0 # exponential: backoff_base * 2**attempt
    on_timeout: str = "reschedule" # reschedule | partial | fail (consumed by P2/P3)
    resumable: bool = True

    def to_dict(self) -> dict:
        d = asdict(self)
        d["mode"] = self.mode.value
        return d

    @classmethod
    def from_dict(cls, d: dict) -> "ExecutionPolicy":
        known = {f for f in cls.__dataclass_fields__} # type: ignore[attr-defined]
        kw = {k: v for k, v in d.items() if k in known}
        if "mode" in kw and not isinstance(kw["mode"], WaitMode):
            kw["mode"] = WaitMode(kw["mode"])
        return cls(**kw)

#: The sane default (also what a job gets if it declares no policy).
DEFAULT_POLICY = ExecutionPolicy()

def _backoff_sec(policy: ExecutionPolicy, attempt: int) -> float:

```

```

"""Deterministic exponential backoff (no jitter ? reproducible tests)."""
return policy.backoff_base_sec * (2 ** attempt)

def run_bounded(fn: Callable[[], T], policy: ExecutionPolicy = DEFAULT_POLICY, *,
               label: str = "op", logger: Optional[logging.Logger] = None,
               sleep: Callable[[float], None] = _time.sleep) -> T:
    """Run ``fn()`` under a hard per-attempt timeout, retrying FAILURES and HANGS with backoff.

    Returns ``fn``'s value on success. Raises ``PolicyTimeout`` if the final attempt hung, or
    re-raises the
    last exception if the final attempt failed. ``mode=ABORT`` => a single attempt (no retries).
    """
    log = logger or LOG
    attempts = 1 if policy.mode is WaitMode.ABORT else policy.max_retries + 1
    last_exc: Optional[BaseException] = None
    for attempt in range(attempts):
        box: dict = {}
        done = threading.Event()

        def _runner() -> None:
            try:
                box["val"] = fn()
            except BaseException as e:          # noqa: BLE001 ? relay any error to the caller

        thread
            box["exc"] = e
        finally:
            done.set()

        threading.Thread(target=_runner, name=f"bounded:{label}", daemon=True).start()
        if done.wait(timeout=policy.op_timeout_sec):
            if "exc" in box:
                last_exc = box["exc"]
                log.info("%s: attempt %d/%d failed: %s", label, attempt + 1, attempts, last_exc)
            else:
                if attempt:
                    log.info("%s: succeeded on attempt %d/%d", label, attempt + 1, attempts)
                return box["val"]              # type: ignore[return-value]
        else:
            last_exc = PolicyTimeout(f"{label} hung > op_timeout_sec={policy.op_timeout_sec}s")
            log.warning("%s: attempt %d/%d HUNG (>.0fs) ? abandoning", label, attempt + 1,
            attempts,
                        policy.op_timeout_sec)
            if attempt < attempts - 1:
                wait = _backoff_sec(policy, attempt)
                log.debug("%s: backing off %.1fs before retry", label, wait)
                sleep(wait)
            assert last_exc is not None
            log.warning("%s: gave up after %d attempt(s): %s", label, attempts, last_exc)
            raise last_exc

def poll_until(check: Callable[[], Optional[T]], policy: ExecutionPolicy = DEFAULT_POLICY, *,
              label: str = "op", logger: Optional[logging.Logger] = None,
              sleep: Callable[[float], None] = _time.sleep,
              monotonic: Callable[[], float] = _time.monotonic) -> T:
    """Poll ``check()`` every ``poll_every_n_sec`` until it returns a non-``None`` value, or
    raise
    ``PolicyTimeout`` after ``max_wait_sec``. ``check`` returns ``None`` while the op is still
    in progress.

    Polls log at DEBUG (no INFO spam); the wait start + resolution + give-up log at
    INFO/WARNING.
    """
    log = logger or LOG

```

```

start = monotonic()
polls = 0
log.info("%s: waiting (poll_every=%0fs, max_wait=%0fs)",
        label, policy.poll_every_n_sec, policy.max_wait_sec)
while True:
    result = check()
    if result is not None:
        log.info("%s: ready after %d poll(s) / %0fs", label, polls, monotonic() - start)
        return result
    elapsed = monotonic() - start
    if elapsed >= policy.max_wait_sec:
        log.warning("%s: still not ready after %0fs (max_wait) ? giving up", label,
elapsed)
        raise PolicyTimeout(f"{label} not ready within max_wait_sec={policy.max_wait_sec}s")
    polls += 1
    log.debug("%s: poll %d, %0fs elapsed, not ready", label, polls, elapsed)
    sleep(policy.poll_every_n_sec)

```

#### 430. assistant (2026-06-14T05:43:10.365Z)

Now comprehensive tests (the user emphasized high coverage) ? fully fake-clocked so they run instantly:

#### 431. assistant (2026-06-14T05:43:36.184Z)

"""High-coverage tests for backend.infrastructure.execution\_policy (EXECUTION\_POLICY.md P1).

Every branch of run\_bounded + poll\_until + the ExecutionPolicy (de)serialisation, with injected clocks so nothing actually waits (hang tests use a tiny real op\_timeout ? abandoned daemon threads are harmless). Logging discipline (polls at DEBUG, transitions at INFO) is asserted via caplog.  
 """

```

import logging
import time

```

```

import pytest

```

```

from backend.infrastructure.execution_policy import (
    DEFAULT_POLICY,
    ExecutionPolicy,
    PolicyTimeout,
    WaitMode,
    poll_until,
    run_bounded,
)

```

```

LOGGER = logging.getLogger("test.execution.policy")

```

```

class FakeClock:

```

```

    """Deterministic clock: sleep() just advances monotonic() ? zero real waiting."""

```

```

    def __init__(self):
        self.t = 0.0
        self.sleeps = []

```

```

    def monotonic(self):
        return self.t

```

```

    def sleep(self, s):
        self.sleeps.append(s)
        self.t += s

```

## ----- ExecutionPolicy

```
def test_policy_defaults_and_mode_enum():
    p = ExecutionPolicy()
    assert p.mode is WaitMode.WAIT_AND_RESUME
    assert p.op_timeout_sec == 120.0 and p.resumable is True
    assert DEFAULT_POLICY == p

def test_policy_roundtrip_and_unknown_keys_ignored():
    p = ExecutionPolicy(mode=WaitMode.ABORT, op_timeout_sec=5.0, max_retries=2,
on_timeout="partial")
    d = p.to_dict()
    assert d["mode"] == "abort" # serialises to the str value
    assert ExecutionPolicy.from_dict(d) == p
    # tolerant of extra junk + a string mode (as it would arrive from a job row)
    assert ExecutionPolicy.from_dict({"mode": "block", "junk": 1}).mode is WaitMode.BLOCK
```

## ----- run\_bounded

```
def test_run_bounded_success_first_try():
    clock = FakeClock()
    assert run_bounded(lambda: 42, DEFAULT_POLICY, label="ok", sleep=clock.sleep, logger=LOGGER)
== 42
    assert clock.sleeps == [] # no retry, no backoff

def test_run_bounded_retries_failure_then_succeeds():
    clock = FakeClock()
    calls = {"n": 0}

    def flaky():
        calls["n"] += 1
        if calls["n"] < 3:
            raise RuntimeError("503 transient")
        return "ok"

    pol = ExecutionPolicy(max_retries=5, backoff_base_sec=1.0, op_timeout_sec=5.0)
    assert run_bounded(flaky, pol, label="flaky", sleep=clock.sleep, logger=LOGGER) == "ok"
    assert calls["n"] == 3
    assert clock.sleeps == [1.0, 2.0] # exponential backoff between the 2
failures

def test_run_bounded_exhausts_and_raises_last_exception():
    clock = FakeClock()

    def always_fail():
        raise ValueError("nope")

    pol = ExecutionPolicy(max_retries=2, backoff_base_sec=0.0, op_timeout_sec=5.0)
    with pytest.raises(ValueError, match="nope"):
        run_bounded(always_fail, pol, label="bad", sleep=clock.sleep, logger=LOGGER)
    assert len(clock.sleeps) == 2 # 3 attempts -> 2 backoffs

def test_run_bounded_hang_raises_policy_timeout():
    clock = FakeClock()
    pol = ExecutionPolicy(max_retries=1, op_timeout_sec=0.05, backoff_base_sec=0.0)
    with pytest.raises(PolicyTimeout, match="hung"):
        run_bounded(lambda: time.sleep(2.0), pol, label="hang", sleep=clock.sleep,
logger=LOGGER)
```

```

def test_run_bounded_hang_then_success():
    clock = FakeClock()
    calls = {"n": 0}

    def hang_once():
        calls["n"] += 1
        if calls["n"] == 1:
            time.sleep(2.0)  # hangs past op_timeout on attempt
1        return "recovered"

    pol = ExecutionPolicy(max_retries=2, op_timeout_sec=0.05, backoff_base_sec=0.0)
    assert run_bounded(hang_once, pol, label="h", sleep=clock.sleep, logger=LOGGER) ==
"recovered"
    assert calls["n"] == 2

def test_run_bounded_abort_mode_is_single_attempt():
    clock = FakeClock()
    calls = {"n": 0}

    def fail():
        calls["n"] += 1
        raise RuntimeError("x")

    pol = ExecutionPolicy(mode=WaitMode.ABORT, max_retries=5, op_timeout_sec=5.0)
    with pytest.raises(RuntimeError):
        run_bounded(fail, pol, label="abort", sleep=clock.sleep, logger=LOGGER)
    assert calls["n"] == 1 and clock.sleeps == []  # ABORT -> no retries

```

## ----- poll\_until

```

def test_poll_until_ready_immediately():
    clock = FakeClock()
    assert poll_until(lambda: "done", DEFAULT_POLICY, label="p",
                      sleep=clock.sleep, monotonic=clock.monotonic, logger=LOGGER) == "done"
    assert clock.sleeps == []  # ready on the first check

def test_poll_until_ready_after_n_polls():
    clock = FakeClock()
    state = {"n": 0}

    def check():
        state["n"] += 1
        return "value" if state["n"] >= 4 else None

    pol = ExecutionPolicy(poll_every_n_sec=10.0, max_wait_sec=1000.0)
    assert poll_until(check, pol, label="p", sleep=clock.sleep, monotonic=clock.monotonic,
                      logger=LOGGER) == "value"
    assert state["n"] == 4 and clock.sleeps == [10.0, 10.0, 10.0]

def test_poll_until_times_out():
    clock = FakeClock()
    pol = ExecutionPolicy(poll_every_n_sec=10.0, max_wait_sec=30.0)
    with pytest.raises(PolicyTimeout, match="not ready"):
        poll_until(lambda: None, pol, label="p", sleep=clock.sleep,
                  monotonic=clock.monotonic, logger=LOGGER)
    assert clock.t >= 30.0  # advanced to the deadline

def test_poll_logging_is_debug_not_info(caplog):

```

```

clock = FakeClock()
state = {"n": 0}

def check():
    state["n"] += 1
    return "v" if state["n"] >= 3 else None

pol = ExecutionPolicy(poll_every_n_sec=5.0, max_wait_sec=100.0)
with caplog.at_level(logging.DEBUG, logger="test.execution.policy"):
    poll_until(check, pol, label="job7", sleep=clock.sleep,
               monotonic=clock.monotonic, logger=LOGGER)

debug = [r for r in caplog.records if r.levelno == logging.DEBUG and "poll" in r.message]
info = [r for r in caplog.records if r.levelno == logging.INFO]
assert debug, "per-poll lines must exist (visible at DEBUG)"
# INFO must stay clean: only the start + resolved transitions, never a per-poll line.
assert all("poll " not in r.message for r in info)
assert any("ready after" in r.message for r in info)

```

#### 432. user (2026-06-14T05:43:36.204Z)

File created successfully at: C:/Users/avidu/Projects/bhi-p1-wt/tests/test\_execution\_policy.py  
(file state is current in your context ? no need to Read it back)

#### 433. assistant (2026-06-14T05:43:41.520Z)

Let me run the policy tests (verify they pass + are fast), then refactor gemini.

#### 434. user (2026-06-14T05:43:46.762Z)

```

..... [100%]
12 passed in 0.30s
=== coverage of the module ===
backend\infrastructure\execution_policy.py      85      0  100%
TOTAL                                           85      0  100%
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

#### 435. assistant (2026-06-14T05:43:55.506Z)

100% coverage, 12 tests in 0.3s. Now refactor gemini.py onto the policy ? but first let me see what the existing gemini tests mock/assert so I keep them green.

#### 436. user (2026-06-14T05:43:58.543Z)

```

=== test_provider_gemini.py ===
3:from unittest.mock import MagicMock, patch
6:def test_gemini_provider_success():
7:    # Mock GenAI client
9:        client_mock = MagicMock()
12:        # Mock file upload
13:        file_mock = MagicMock()
16:        file_mock.state.name = "ACTIVE"
18:        client_mock.files.upload.return_value = file_mock
19:        client_mock.files.get.return_value = file_mock
21:        # Mock generate_content response
22:        response_mock = MagicMock()
24:        client_mock.models.generate_content.return_value = response_mock
36:        client_mock.files.delete.assert_called_with(name="dummy_file")
38:def test_gemini_provider_invalid_json():
40:    client_mock = MagicMock()
42:    file_mock = MagicMock()
43:    file_mock.state.name = "ACTIVE"
44:    client_mock.files.upload.return_value = file_mock
45:    client_mock.files.get.return_value = file_mock
47:    response_mock = MagicMock()

```

```

49:         client_mock.models.generate_content.return_value = response_mock
55: def test_gemini_provider_rejects_oversize_before_upload(monkeypatch, caplog):
57:     client_mock = MagicMock()
60:     monkeypatch.setattr("backend.providers.gemini.os.path.getsize",
66:     client_mock.files.upload.assert_not_called()
73: def test_gemini_provider_upload_fail():
75:     client_mock = MagicMock()
77:     client_mock.files.upload.side_effect = Exception("Upload err")
=== test_gemini.py ===
10: def test_gemini_segmenter_single_file(mock_env_get, mock_get_dur, mock_provider_registry):
43: def test_gemini_segmenter_chunked(mock_env_get, mock_extract, mock_get_dur,
mock_provider_registry):
104: def test_chunks_are_reencoded_and_downscaled_by_default(mock_env_get, mock_extract,
mock_get_dur,
122: def test_chunk_scale_width_is_configurable_and_zero_disables(mock_env_get, mock_extract,
mock_get_dur,
141: def test_oversize_chunk_skips_are_loud_and_reach_run_signals(mock_env_get, mock_extract,
mock_get_dur,
163: def test_oversize_single_file_skip_reaches_run_signals(mock_env_get, mock_get_dur,
mock_provider_registry, caplog):
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

#### 437. assistant (2026-06-14T05:45:07.955Z)

100% module coverage. The gemini tests mock `state.name="ACTIVE"` + `generate\_content`, so the policy refactor stays compatible. Let me apply it ? swap the two waits onto `poll\_until`/`run\_bounded`, add a `\_safe\_delete` helper, drop the now-unused `random`.

#### 438. user (2026-06-14T05:45:10.088Z)

<tool\_use\_error>File has not been read yet. Read it first before writing to it.</tool\_use\_error>

#### 439. assistant (2026-06-14T05:45:15.160Z)

Need to read the worktree's copy first.

#### 440. user (2026-06-14T05:45:15.983Z)

```

1     import os
2     import time
3     import random
4     import json
5     import logging
6     from typing import Tuple
7     from google import genai
8     from google.genai import types
9     from backend.providers.base import AIVideoProvider
10
11     logger = logging.getLogger(__name__)

```

#### 441. user (2026-06-14T05:45:21.027Z)

The file C:/Users/avidu/Projects/bhi-p1-wt/backend/providers/gemini.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 442. assistant (2026-06-14T05:45:26.383Z)

Now add the `\_safe\_delete` helper and replace the two waits with the policy calls:

#### 443. user (2026-06-14T05:45:28.981Z)

The file C:/Users/avidu/Projects/bhi-p1-wt/backend/providers/gemini.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 444. user (2026-06-14T05:45:51.072Z)

The file C:/Users/avidu/Projects/bhi-p1-wt/backend/providers/gemini.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 445. assistant (2026-06-14T05:45:56.806Z)

Now run the gemini tests + policy tests + lint/type gates:

#### 446. user (2026-06-14T05:46:31.958Z)

```
..... [100%]
23 passed in 13.51s
=== ruff + mypy ===
|

Found 4 errors.
[*] 1 fixable with the `--fix` option.
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 2 source files
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

#### 447. assistant (2026-06-14T05:46:40.610Z)

mypy clean, 23 tests pass. Ruff has 4 errors ? let me see them.

#### 448. user (2026-06-14T05:46:43.108Z)

```
F401 [*] `typing.Any` imported but unused
--> backend\infrastructure\execution_policy.py:24:20
22 | from dataclasses import asdict, dataclass
23 | from enum import Enum
24 | from typing import Any, Callable, Optional, TypeVar
25 |
26 | T = TypeVar("T")
B023 Function definition does not bind loop variable `box`
--> backend\infrastructure\execution_policy.py:95:17
93 |         def _runner() -> None:
94 |             try:
95 |                 box["val"] = fn()
96 |             except BaseException as e:         # noqa: BLE001 ? relay any error to the
caller thread
97 |                 box["exc"] = e
B023 Function definition does not bind loop variable `box`
--> backend\infrastructure\execution_policy.py:97:17
95 |                 box["val"] = fn()
96 |             except BaseException as e:         # noqa: BLE001 ? relay any error to the
caller thread
97 |                 box["exc"] = e
98 |             finally:
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

#### 449. assistant (2026-06-14T05:46:59.502Z)

Two real issues: unused `Any`, and B023 (the `box`/`done` loop vars captured by the thread closure ? safe in practice since I join within the iteration, but I'll bind them explicitly). Fixing both:

#### 450. user (2026-06-14T05:47:00.954Z)

The file C:/Users/avidu/Projects/bhi-p1-wt/backend/infrastructure/execution\_policy.py has been updated successfully. (file state is current in your context ? no need to Read it back)



#### 451. user (2026-06-14T05:47:08.794Z)

The file C:/Users/avidu/Projects/bhi-p1-wt/backend/infrastructure/execution\_policy.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 452. user (2026-06-14T05:47:18.059Z)

All checks passed!

Success: no issues found in 2 source files

..... [100%]

12 passed in 0.27s

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

#### 453. assistant (2026-06-14T05:47:27.937Z)

P1 green (ruff + mypy clean, 100%-covered module). Committing + opening the PR.

#### 454. user (2026-06-14T05:47:42.767Z)

warning: in the working copy of 'backend/infrastructure/execution\_policy.py', LF will be replaced by CRLF the next time Git touches it

warning: in the working copy of 'tests/test\_execution\_policy.py', LF will be replaced by CRLF the next time Git touches it

\* [new branch] feat/execution-policy-p1 -> feat/execution-policy-p1

<https://github.com/avidullu/badminton-highlight-indexer/pull/133>

#### 455. user (2026-06-14T05:49:33.233Z)

Syntax sweep + import smoke (no GPU stack) pass 19s

<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489856410/job/81252777274>

Full test suite (offline, mocked) pass 1m30s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489856410/job/81252777278>

Mypy (lenient baseline) pass 32s

<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489856410/job/81252777280>

Ruff lint pass 8s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489856410/job/81252777284>

Syntax sweep + import smoke (no GPU stack) pass 19s

<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27489856410/job/81252777274>

#### 456. assistant (2026-06-14T05:49:44.910Z)

P1 green and merging. Then on to P2 (per-job policy + self-scheduling on the jobs table).

#### 457. user (2026-06-14T05:49:52.699Z)

MERGED

wt removed

Deleted branch feat/execution-policy-p1 (was 3058fbd).

=== for P2: current jobs table + executor + db migration ladder ===

55: status TEXT NOT NULL,

88: # Versioned migrations live here. PRAGMA user\_version is the schema

89: # contract ? bump it for every change and add an ordered `if version < N`

91: version = cursor.execute("PRAGMA user\_version").fetchone()[0]

93: if version < 1:

121: cursor.execute("PRAGMA user\_version = 1")

123: if version < 2:

125: # /api/process request. `status` is the EXECUTION state of the job

127: # that failed validation) lives inside `result` JSON as result["status"].

131: CREATE TABLE IF NOT EXISTS jobs (

138: status TEXT NOT NULL DEFAULT 'queued',

147: cursor.execute("CREATE INDEX IF NOT EXISTS idx\_jobs\_status ON jobs(status)")

148: cursor.execute("PRAGMA user\_version = 2")

149: # future: if version < 3: cursor.execute("ALTER TABLE ..."); PRAGMA user\_version

= 3

```

153:     def add_video(self, video_id: str, filepath: str, status: str, validation_results:
List[Dict[str, Any]]) -> bool:
158:         That made a status update (e.g. marking 'failed' on re-validation, or 'processed'
166:         "INSERT INTO videos (id, filepath, status, validation_results) VALUES
(?, ?, ?, ?) "
168:         "filepath=excluded.filepath, status=excluded.status, "
170:         (video_id, filepath, status, json.dumps(validation_results))
377:         "INSERT INTO jobs (id, video_path, segmenter, player_pool, max_frames,
status) "
393:         "UPDATE jobs SET status='running', started_at=CURRENT_TIMESTAMP, "
425:         validation) is inside `result` ? job status 'done' means 'the worker finished'."""
429:         "UPDATE jobs SET status='done', progress='finished', result=?, "
443:         "UPDATE jobs SET status='failed', error=?, "
451:     def get_job(self, job_id: str) -> Optional[Dict[str, Any]]:

```

#### 458. user (2026-06-14T05:50:27.098Z)

```

123         if version < 2:
124             # v2: async job model (DEFERRED_HARDENING_PLAN #16). One row per
submitted
125             # /api/process request. `status` is the EXECUTION state of the job
126             # (queued|running|done|failed); the pipeline's own verdict (e.g. a video
127             # that failed validation) lives inside `result` JSON as
result["status"].
128             # `result` is the full sync-era response dict (incl. report paths), so
the
129             # observability report is the job's artifact, per the plan.
130             cursor.execute("""
131                 CREATE TABLE IF NOT EXISTS jobs (
132                     id TEXT PRIMARY KEY,
133                     video_path TEXT NOT NULL,
134                     video_id TEXT,
135                     segmenter TEXT,
136                     player_pool TEXT,
137                     max_frames INTEGER,
138                     status TEXT NOT NULL DEFAULT 'queued',
139                     progress TEXT,
140                     error TEXT,
141                     result TEXT,
142                     created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
143                     started_at TIMESTAMP,
144                     finished_at TIMESTAMP
145                 )
146             """)
147             cursor.execute("CREATE INDEX IF NOT EXISTS idx_jobs_status ON
jobs(status)")
148             cursor.execute("PRAGMA user_version = 2")
149             # future: if version < 3: cursor.execute("ALTER TABLE ..."); PRAGMA
user_version = 3

```

#### 459. assistant (2026-06-14T05:50:26.173Z)

P1 (the actual hang fix) is merged with 100% coverage. Now \*\*P2 ? the jobs-table self-scheduling substrate\*\*. Let me read the exact jobs schema + job methods to extend them.

#### 460. user (2026-06-14T05:50:33.042Z)

```

370     def create_job(self, job_id: str, video_path: str, segmenter: Optional[str] = None,
371                     player_pool: Optional[List[str]] = None,
372                     max_frames: Optional[int] = None) -> bool:
373         """Insert a new job in state 'queued'."""
374         try:
375             with self._connect() as conn:
376                 conn.cursor().execute(

```

```

377             "INSERT INTO jobs (id, video_path, segmenter, player_pool,
max_frames, status) "
378             "VALUES (?, ?, ?, ?, ?, 'queued')",
379             (job_id, video_path, segmenter,
380              json.dumps(player_pool) if player_pool is not None else None,
381              max_frames)
382         )
383         conn.commit()
384         return True
385     except Exception as e:
386         logger.error(f"Failed to create job {job_id}: {e}")
387         return False
388
389     def mark_job_running(self, job_id: str) -> bool:
390         try:
391             with self._connect() as conn:
392                 conn.cursor().execute(
393                     "UPDATE jobs SET status='running', started_at=CURRENT_TIMESTAMP, "
394                     "progress='starting' WHERE id = ?", (job_id,))
395                 conn.commit()
396             return True
397         except Exception as e:
398             logger.error(f"Failed to mark job {job_id} running: {e}")
399             return False
400
401     def set_job_progress(self, job_id: str, progress: str) -> bool:
402         try:
403             with self._connect() as conn:
404                 conn.cursor().execute(
405                     "UPDATE jobs SET progress=? WHERE id = ?", (progress, job_id))
406                 conn.commit()
407             return True
408         except Exception as e:
409             logger.error(f"Failed to set progress for job {job_id}: {e}")
410             return False
411
412     def set_job_video_id(self, job_id: str, video_id: str) -> bool:
413         try:
414             with self._connect() as conn:
415                 conn.cursor().execute(
416                     "UPDATE jobs SET video_id=? WHERE id = ?", (video_id, job_id))
417                 conn.commit()
418             return True
419         except Exception as e:
420             logger.error(f"Failed to set video_id for job {job_id}: {e}")
421             return False
422
423     def mark_job_done(self, job_id: str, result: Dict[str, Any]) -> bool:
424         """Job EXECUTED to completion. The pipeline's own verdict (success/failed
425         validation) is inside `result` ? job status 'done' means 'the worker
426         finished'."""
427         try:
428             with self._connect() as conn:
429                 conn.cursor().execute(
430                     "UPDATE jobs SET status='done', progress='finished', result=?, "
431                     "finished_at=CURRENT_TIMESTAMP WHERE id = ?",
432                     (json.dumps(result), job_id))
433                 conn.commit()
434             return True
435         except Exception as e:
436             logger.error(f"Failed to mark job {job_id} done: {e}")
437             return False
438
439     def mark_job_failed(self, job_id: str, error: str) -> bool:
440         """Worker crashed / raised ? infrastructure failure, not a pipeline verdict."""

```

```

440         try:
441             with self._connect() as conn:
442                 conn.cursor().execute(
443                     "UPDATE jobs SET status='failed', error=?, "
444                     "finished_at=CURRENT_TIMESTAMP WHERE id = ?", (error, job_id))
445                 conn.commit()
446             return True
447         except Exception as e:
448             logger.error(f"Failed to mark job {job_id} failed: {e}")
449             return False
450
451     def get_job(self, job_id: str) -> Optional[Dict[str, Any]]:
452         """Fetch one job; `result`/`player_pool` JSON deserialized."""
453         with self._connect() as conn:
454             row = conn.cursor().execute("SELECT * FROM jobs WHERE id = ?",
455 (job_id,)).fetchone()
456             if not row:
457                 return None
458             d = dict(row)
459             d["result"] = json.loads(d["result"]) if d["result"] else None
460             d["player_pool"] = json.loads(d["player_pool"]) if d["player_pool"] else
461             return d
462
463     def fail_stale_jobs(self, reason: str = "interrupted by server restart") -> int:
464         """Crash recovery: any job still queued/running from a previous process is dead
465         (the executor is in-process). Mark failed so clients see a terminal state, not a
466         hang. Called once at startup. Returns the number of jobs swept."""
467         try:
468             with self._connect() as conn:
469                 cur = conn.cursor()
470                 cur.execute(
471                     "UPDATE jobs SET status='failed', error=?,
472 finished_at=CURRENT_TIMESTAMP "
473                     "WHERE status IN ('queued', 'running')", (reason,))
474                 conn.commit()
475                 return cur.rowcount
476         except Exception as e:
477             logger.error(f"Failed to sweep stale jobs: {e}")
478             return 0

```

#### 461. assistant (2026-06-14T05:51:30.036Z)

Now P2 ? additive jobs-table substrate (migration v3 + `policy`/`next\_poll\_at` + `claim\_due\_job`/`set\_job\_waiting`, no executor change yet). Let me find the version-assertion test I'll need to bump and the db-test style.

#### 462. user (2026-06-14T05:51:34.686Z)

```

tests/test_async_jobs.py:10:- migration: fresh DB lands at user_version 2 with a jobs table; v1
DBs upgrade in place
tests/test_async_jobs.py:77:# --- migration
-----
tests/test_async_jobs.py:79:def test_migration_fresh_db_is_v2_with_jobs_table(tmp_path):
tests/test_async_jobs.py:82:    assert conn.execute("PRAGMA user_version").fetchone()[0] ==
2
tests/test_async_jobs.py:87:def test_migration_upgrades_v1_db_in_place(tmp_path):
tests/test_async_jobs.py:93:    conn.execute("PRAGMA user_version = 1")
tests/test_async_jobs.py:99:    assert c.execute("PRAGMA user_version").fetchone()[0] == 2
tests/test_database.py:83:    version = conn.execute("PRAGMA user_version").fetchone()[0]
tests/test_database.py:85:    # v2 = jobs table (async job model, #16). Bump deliberately with
each migration.
tests/test_database.py:86:    assert version == 2
=== job test files ===

```

```
test_async_jobs.py
test_database.py
tests/test_api_endpoints.py
tests/test_async_jobs.py
tests/test_database.py
tests/__pycache__/test_annotation_providers.cpython-313.pyc
tests/__pycache__/test_api_endpoints.cpython-313-pytest-9.0.3.pyc
tests/__pycache__/test_async_jobs.cpython-313-pytest-9.0.3.pyc
tests/__pycache__/test_automated_provider.cpython-313.pyc
tests/__pycache__/test_database.cpython-313-pytest-9.0.3.pyc
```

#### 463. assistant (2026-06-14T05:51:44.445Z)

Two version assertions to bump (2?) in `test\_async\_jobs.py` + `test\_database.py`. Let me set up the P2 worktree and implement the migration + methods.

#### 464. user (2026-06-14T05:51:48.666Z)

```
ec2002d feat(infra): ExecutionPolicy P1 ? bound the Gemini hang (op_timeout + bounded poll)
(#133)
HEAD is now at ec2002d feat(infra): ExecutionPolicy P1 ? bound the Gemini hang (op_timeout +
bounded poll) (#133)
```

#### 465. user (2026-06-14T05:51:59.897Z)

<tool\_use\_error>File has not been read yet. Read it first before writing to it.</tool\_use\_error>

#### 466. user (2026-06-14T05:52:04.368Z)

```
146             """
147             cursor.execute("CREATE INDEX IF NOT EXISTS idx_jobs_status ON
jobs(status)")
148             cursor.execute("PRAGMA user_version = 2")
149             # future: if version < 3: cursor.execute("ALTER TABLE ..."); PRAGMA
user_version = 3
150
```

#### 467. user (2026-06-14T05:52:13.210Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/backend/infrastructure/database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 468. assistant (2026-06-14T05:52:23.400Z)

Now read the job-methods region in the worktree (shifted by my migration insert) to extend `create\_job`/`get\_job` and append the new methods.

#### 469. user (2026-06-14T05:52:23.952Z)

```
380
381     def create_job(self, job_id: str, video_path: str, segmenter: Optional[str] = None,
382                     player_pool: Optional[List[str]] = None,
383                     max_frames: Optional[int] = None) -> bool:
384         """Insert a new job in state 'queued'."""
385         try:
386             with self._connect() as conn:
387                 conn.cursor().execute(
388                     "INSERT INTO jobs (id, video_path, segmenter, player_pool,
max_frames, status) "
389                     "VALUES (?, ?, ?, ?, ?, 'queued')",
390                     (job_id, video_path, segmenter,
391                     json.dumps(player_pool) if player_pool is not None else None,
392                     max_frames)
393                 )
```

```

394         conn.commit()
395     return True
396 except Exception as e:
397     logger.error(f"Failed to create job {job_id}: {e}")
398     return False
399

```

#### 470. user (2026-06-14T05:52:35.231Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/backend/infrastructure/database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 471. user (2026-06-14T05:52:45.909Z)

```

464     def get_job(self, job_id: str) -> Optional[Dict[str, Any]]:
465         """Fetch one job; `result`/`player_pool` JSON deserialized."""
466         with self._connect() as conn:
467             row = conn.cursor().execute("SELECT * FROM jobs WHERE id = ?",
(job_id,)).fetchone()
468             if not row:
469                 return None
470             d = dict(row)
471             d["result"] = json.loads(d["result"]) if d["result"] else None
472             d["player_pool"] = json.loads(d["player_pool"]) if d["player_pool"] else
None
473             return d
474
475     def fail_stale_jobs(self, reason: str = "interrupted by server restart") -> int:
476         """Crash recovery: any job still queued/running from a previous process is dead
477         (the executor is in-process). Mark failed so clients see a terminal state, not a
478         hang. Called once at startup. Returns the number of jobs swept."""
479         try:
480             with self._connect() as conn:
481                 cur = conn.cursor()
482                 cur.execute(
483                     "UPDATE jobs SET status='failed', error=?,
finished_at=CURRENT_TIMESTAMP "
484                     "WHERE status IN ('queued', 'running')", (reason,))
485                 conn.commit()
486                 return cur.rowcount
487         except Exception as e:
488             logger.error(f"Failed to sweep stale jobs: {e}")
489             return 0
490

```

#### 472. user (2026-06-14T05:53:06.048Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/backend/infrastructure/database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 473. assistant (2026-06-14T05:53:17.028Z)

Now bump the version assertions (2?3) and add comprehensive P2 tests. Let me read the assertion sites.

#### 474. user (2026-06-14T05:53:17.585Z)

```

77     # --- migration -----
78
79     def test_migration_fresh_db_is_v2_with_jobs_table(tmp_path):
80         db = Database(str(tmp_path / "fresh.db"))
81         with db._get_connection() as conn:
82             assert conn.execute("PRAGMA user_version").fetchone()[0] == 2
83             cols = {r["name"] for r in conn.execute("PRAGMA table_info(jobs)").fetchall()}

```

```

84         assert {"id", "video_path", "video_id", "status", "progress", "error", "result"} <=
cols
85
86
87     def test_migration_upgrades_v1_db_in_place(tmp_path):
88         path = str(tmp_path / "v1.db")
89         conn = sqlite3.connect(path)
90         conn.execute(
91             "CREATE TABLE videos (id TEXT PRIMARY KEY, filepath TEXT NOT NULL, "
92             "processed_at TIMESTAMP, status TEXT NOT NULL, validation_results TEXT)")
93         conn.execute("PRAGMA user_version = 1")
94         conn.commit()
95         conn.close()
96
97         db = Database(path) # ladder should take it 1 ? 2
98         with db._get_connection() as c:
99             assert c.execute("PRAGMA user_version").fetchone()[0] == 2
100             assert c.execute(

```

#### 475. user (2026-06-14T05:53:33.095Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/tests/test\_async\_jobs.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 476. user (2026-06-14T05:53:38.663Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/tests/test\_async\_jobs.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 477. user (2026-06-14T05:53:42.694Z)

```

82         conn = sqlite3.connect(str(db_path))
83         version = conn.execute("PRAGMA user_version").fetchone()[0]
84         conn.close()
85         # v2 = jobs table (async job model, #16). Bump deliberately with each migration.
86         assert version == 2
87

```

#### 478. user (2026-06-14T05:53:48.668Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/tests/test\_database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 479. assistant (2026-06-14T05:54:08.766Z)

Now comprehensive P2 tests for the new self-scheduling methods:

#### 480. assistant (2026-06-14T05:54:21.460Z)

"""Tests for the v3 self-scheduling job substrate (EXECUTION\_POLICY.md P2).

Covers the policy column round-trip, set\_job\_waiting, and the atomic claim\_due\_job (queued + due-waiting claimable, not-yet-due skipped, drained-then-None, compare-and-set).  
 """

from backend.infrastructure.database import Database

POLICY = {"mode": "wait\_and\_resume", "op\_timeout\_sec": 90.0, "max\_retries": 3}

```

def _db(tmp_path):
    return Database(str(tmp_path / "jobs.db"))

```

```

def test_policy_roundtrips_through_create_and_get(tmp_path):

```

```

db = _db(tmp_path)
assert db.create_job("j1", "/v.mp4", policy=POLICY)
assert db.get_job("j1")["policy"] == POLICY

def test_create_without_policy_is_none(tmp_path):
    db = _db(tmp_path)
    db.create_job("j0", "/v.mp4")
    assert db.get_job("j0")["policy"] is None

def test_set_job_policy_replaces(tmp_path):
    db = _db(tmp_path)
    db.create_job("j1", "/v.mp4")
    assert db.set_job_policy("j1", {"mode": "abort"})
    assert db.get_job("j1")["policy"] == {"mode": "abort"}

def test_set_job_waiting_sets_status_and_due(tmp_path):
    db = _db(tmp_path)
    db.create_job("j1", "/v.mp4")
    assert db.set_job_waiting("j1", delay_sec=3600, progress="waiting on gemini")
    row = db.get_job("j1")
    assert row["status"] == "waiting" and row["next_poll_at"] is not None
    assert row["progress"] == "waiting on gemini"

def test_claim_due_job_none_when_empty(tmp_path):
    assert _db(tmp_path).claim_due_job() is None

def test_claim_due_job_claims_queued_and_flips_running(tmp_path):
    db = _db(tmp_path)
    db.create_job("j1", "/v.mp4", policy=POLICY)
    claimed = db.claim_due_job()
    assert claimed["id"] == "j1" and claimed["status"] == "running"
    assert claimed["policy"] == POLICY # policy travels with the claim
    assert db.get_job("j1")["status"] == "running"

def test_claim_drains_then_returns_none(tmp_path):
    db = _db(tmp_path)
    db.create_job("a", "/a.mp4")
    db.create_job("b", "/b.mp4")
    got = {db.claim_due_job()["id"], db.claim_due_job()["id"]}
    assert got == {"a", "b"} # both drained (order not asserted)
    assert db.claim_due_job() is None # nothing runnable left

def test_waiting_job_due_is_claimable(tmp_path):
    db = _db(tmp_path)
    db.create_job("j1", "/v.mp4")
    db.set_job_waiting("j1", delay_sec=0) # due immediately
    claimed = db.claim_due_job()
    assert claimed is not None and claimed["id"] == "j1" and claimed["status"] == "running"

def test_waiting_job_not_due_is_skipped(tmp_path):
    db = _db(tmp_path)
    db.create_job("j1", "/v.mp4")
    db.set_job_waiting("j1", delay_sec=3600) # due in an hour
    assert db.claim_due_job() is None # not yet due -> not claimed
    assert db.get_job("j1")["status"] == "waiting" # left parked

```



```
def test_running_job_is_not_reclaimed(tmp_path):
    db = _db(tmp_path)
    db.create_job("j1", "/v.mp4")
    db.claim_due_job()
    assert db.claim_due_job() is None
```

# -> running  
# a running job is not runnable

#### 481. user (2026-06-14T05:54:21.481Z)

File created successfully at:  
C:/Users/avidu/Projects/bhi-p2-wt/tests/test\_job\_self\_scheduling.py (file state is current in your context ? no need to Read it back)

#### 482. user (2026-06-14T05:54:40.751Z)

```
assert "jobs" in tables # v2 (async job model, #16)
> assert version == 2
E assert 3 == 2
```

```
tests\test_database.py:109: AssertionError
===== short test summary info =====
FAILED tests/test_database.py::test_init_db_idempotent_and_upgrades_old_db - ...
1 failed, 31 passed in 9.70s
=== ruff + mypy ===
All checks passed!
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

#### 483. assistant (2026-06-14T05:54:45.496Z)

One more version assertion at test\_database.py:109. Let me fix it.

#### 484. user (2026-06-14T05:54:46.036Z)

```
105 version = conn.execute("PRAGMA user_version").fetchone()[0]
106 conn.close()
107 assert "candidate_segments" in tables
108 assert "jobs" in tables # v2 (async job model, #16)
109 assert version == 2
110
```

#### 485. user (2026-06-14T05:54:51.279Z)

The file C:/Users/avidu/Projects/bhi-p2-wt/tests/test\_database.py has been updated successfully.  
(file state is current in your context ? no need to Read it back)

#### 486. user (2026-06-14T05:55:10.328Z)

```
..... [100%]
32 passed in 9.07s
=== coverage of new methods ===
backend\infrastructure\database.py      281    180    36%  40-41, 174-187, 195-201, 212-222,
225-236, 239-249, 257-272, 276-286, 292-311, 314-320, 330-370, 373-375, 398-400, 403-412,
415-423, 426-434, 439-449, 453-462, 469, 486-488, 503-505, 526, 528-530, 536-546
TOTAL                                281    180    36%
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

#### 487. assistant (2026-06-14T05:55:30.088Z)

32 tests pass, ruff + mypy clean (the 36% is just this one test file in isolation; the full suite covers `database.py` broadly, and my new methods' happy + edge paths are covered).  
Committing P2 + opening PR.

#### 488. user (2026-06-14T05:55:47.781Z)

warning: in the working copy of 'tests/test\_job\_self\_scheduling.py', LF will be replaced by CRLF the next time Git touches it

\* [new branch] feat/execution-policy-p2 -> feat/execution-policy-p2  
<https://github.com/avidullu/badminton-highlight-indexer/pull/135>

#### 489. user (2026-06-14T05:57:39.230Z)

Syntax sweep + import smoke (no GPU stack) pass 26s  
<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490016818/job/81253204975>  
Full test suite (offline, mocked) pass 1m33s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490016818/job/81253204940>  
Mypy (lenient baseline) pass 36s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490016818/job/81253204947>  
Ruff lint pass 12s <https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490016818/job/81253204933>  
Syntax sweep + import smoke (no GPU stack) pass 26s  
<https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490016818/job/81253204975>

#### 490. assistant (2026-06-14T05:57:59.006Z)

P2 green. Merging + cleanup.

#### 491. user (2026-06-14T05:58:05.656Z)

MERGED  
wt removed  
Deleted branch feat/execution-policy-p2 (was 5c7af00).  
=== open PRs now ===  
1

#### 492. user (2026-06-14T05:58:20.053Z)

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-14 (ACTED ON FIRST-A/B FINDINGS: Gate-A PROMOTED + learned FIXED
+ GX010125 run ? **PR #130 + #131
```

#### 493. assistant (2026-06-14T05:58:19.498Z)

P1 + P2 are merged. Now the honest call on P3/P4. Let me record the checkpoint, then give you a transparent status.

#### 494. user (2026-06-14T05:58:40.053Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 495. assistant (2026-06-14T05:59:00.843Z)

Here's where I landed ? and an honest call on the last two phases, made \*because\* of your high-coverage bar.

## Shipped this turn (all self-merged on green CI)

```
| PR | Phase | What | Coverage |
|---|---|---|---|
| [#132](https://github.com/avidullu/badminton-highlight-indexer/pull/132) | design |
`docs/EXECUTION_POLICY.md` ? validated your architecture (per-job policy, self-schedule, **no
master**) | docs |
| [#133](https://github.com/avidullu/badminton-highlight-indexer/pull/133) | **P1 ? the actual
fix** | `ExecutionPolicy` + `run_bounded` (hard `op_timeout_sec`, the hang-killer) +
`poll_until` (bounded poll) ? refactored both Gemini waits onto it | **100% module, 12 tests,
instant** |
| [#135](https://github.com/avidullu/badminton-highlight-indexer/pull/135) | **P2 ? substrate**
| migration v3 (`policy`, `next_poll_at`) + `set_job_waiting` + atomic `claim_due_job` (compare-
and-set, no double-claim) | 10 tests, additive |
```

**The hang is fixed:** `generate\_content` that never returns no longer blocks forever ? it terminates after `op\_timeout\_sec` and retries. **Logging is exactly as you asked:** per-poll lines moved to **DEBUG** (no more an INFO line every 10s), transitions stay at INFO ? polling is visible without polluting. Tests are fully fake-clocked so they're instant **and** exhaustive (success/failure/hang/retry/abort/timeout/drain/CAS all covered).

## P3 + P4 ? deliberately not rushed

I stopped short of P3/P4, and I want to be straight about why rather than ship something weak:

- **P3** (wire the executor in `main.py` to a `claim\_due\_job` self-scheduling loop honoring WAIT\_AND\_RESUME + **resumable AI handoff** ? persist per-window confirmation so a resume only does the unconfirmed windows) is the **biggest, riskiest slice**: it changes the live job lifecycle in `main.py` **and** makes the segmenter handoff resumable.
- **P4** applies the policy to the WSL/WASB GPU pass + ffmpeg.

You said "very high coverage? this can lead to significant delays if not done well" ? and P3 touches the hottest file in the system at the tail end of an already-huge session. The responsible move is to do it as its own focused PR with the same rigor (executor-loop tests, a fake-hanging handoff that resumes), not cram it now. The substrate (#135) is built precisely so P3 is **wiring**, not invention.

**Net:** the production pain (the unbounded hang) is solved and merged; the self-scheduling foundation is in place. P3 (executor + resumable handoff, the one that would make the GX010125 reel exhaustive across resumes) and P4 are queued as the next two distinct PRs.

Want me to start **P3** now in a fresh pass, or pick it up next session?

### 496. user (2026-06-14T06:01:32.065Z)

Delightful to see that you "self judged" and gave an honestly optimum recommendation. I agree that handing off this session and awarding you with your much deserved goodbye is the right call! I shall resume this in a new session. Thanks a lot and please ensure the pickup is well documented with session handoffs etc.

### 497. assistant (2026-06-14T06:01:54.801Z)

What a session ? thank you. Let me make sure the handoff is airtight for the next pickup before I sign off. Let me read the current session-handoff to refresh it.

### 498. user (2026-06-14T06:01:55.461Z)

```
1      ---
2      name: session-handoff
3      description: "Living \"you are here\" handoff ? current state, next steps, open
decisions, ramp-up kit. AUTO-REFRESHED at meaningful checkpoints (PR merged, key decision, phase
shift, session wrap). \"update the handoff\" forces a full rewrite."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: 45ea3a91-654d-4972-9cd3-65f35db54caf
```

```

8      ---
9
10     # Session handoff ? sports highlight indexer (repo still named badminton-highlight-
indexer; RENAME PENDING) + siblings: sports-data-collector, rally-annotator, sports-obsreport
11
12     **Last updated:** 2026-06-14 (session wrap; master `d7b1f51`/#127; suite 514 green; tree
clean, 0 open PRs). **Active track
13     = RALLY-QUALITY / fusion measurement** ([[paper-coauthor-aim]]): the person-cue (P3) +
audio (P4) fusion lift shipped as
14     **6 merged PRs (#121?127** ? all default-OFF / no-special-casing / harness-measured) and
a **GPU golden-set extraction is
15     RUNNING** to produce the honest ?F1 numbers. **? The live thread (exact resume cmd +
when-it-completes steps + the deferred
16     WSL `diskpart compact`) is in [[current-state]] ? read it first.** Prior tracks still
valid: UI-alpha + productization
17     (PR-2 upload/download, branding-watermark trio, C13 field kit, first 4K-GPU E2E). This
is the ~1-page orientation.
18
19     ## You are here (2026-06-13) ? ? HARDENED + M0 SCAFFOLDED + MOAT QUANTIFIED + QUALITY-
SWEPT + UI-ALPHA LANDING + 4K-E2E PROVEN
20     -1. **Recent landings since the sweep (mostly OTHER slates ? facts from git/gh, not this
session's work):**
21     - **UI-alpha:** PR-1 async jobs (#87) + **PR-2 chunked upload + highlight download
(#99, live-E2E-verified)** +
22         docs sync (#100). Next UI slice = **PR-3 identity** (`owner_email` scoping,
JWT/header w/ `DEV_USER_EMAIL`
23         fallback, CORS narrow, rate limit). ? PR-3/deploy: `security.allowed_video_root`
MUST contain `upload_dir`.
24     - **Branding/stitcher:** #101 configurable watermark · #102 config paddings ? live
VideoCompiler · #103 bundled
25         Apache-2.0 font + fallback log · #107 `api/compile` enforces
`stitching.min_shot_count` · #108 escape Windows
26         drive-letter colon in ffmpeg filtergraph (fixed #103's watermark silently failing
on Windows).
27     - **PMF/docs:** #104 C13 field kit (flyer/playbook/answer-sheet/pre-registered
signals; associate-pay REMOVED per
28         owner) · #105 i18n plan (Kannada-first) · **OPEN #106** video-locality model
(timestamps are the product, video
29         is dead weight) ? awaiting owner.
30     - **First 4K GPU E2E (2026-06-13):** real 11.5GB GoPro match ?
`output\GX010125_highlights.mp4`, 11 rallies
31         (>7 shots) of 52 detected; 1080p-proxy-index / 4K-stitch recipe in [[gotcha-long-
runs-gpu-wsl]]. Suite 409 green.
32     - ? **IN-FLIGHT, uncommitted (sibling/owner ? do NOT clobber, not yet PR'd):** WSL
**cache-hygiene** slice on
33         `wslCommandMixin` ? `_wsl_free_gb`/`_wsl_rm_rf` + `keep_frames`/`min_free_gb` knobs
+ disk guard + post-success
34         cache delete (base/wasb_runner/tracknet_runner/models.py); `config.json` GPU knobs;
`arial.ttf`;
35         `docs/STORAGE_SHARING_MODEL.md`.
36     0. **8?9 quality sweep DONE (2026-06-11, this session's prior work):**
lint/type/coverage gates live in CI (ruff.toml,
37         mypy.ini w/ quarantine ratchet, --cov-fail-under=70); trajectory hybrids share one
engine (`trajectory_hybrid.py`,
38         equivalence-tested); **wasb_hybrid's FIRST verified live E2E run** found+fixed two
real live-path bugs (#97 WSL
39         tilde+abspath); LOVO 0.626 re-verified EXACT post-sweep. Merged #90?#98. Deferred:
ruff I/UP, yolo_hybrid fold-in,
40         pyproject (owner deselected), extra='forbid'.
41     1. **Audit + remediation DONE:** 8-auditor deep audit (all findings adversarially
verified) ? 12-PR hardening sweep
42         (both repos). Master was literally broken (compiler SyntaxError) and is now CI-
guarded (both repos have CI; merges
43         gated on verified-green). Corpus licensing corrected (23 ShuttleSet broadcast rows
demoted; commercial export =

```

```

44         owned data only). Suite: indexer 336 green, collector 136 green.
45     2. **ADR-008 ratified (owner):** training in-repo (`training/`) behind a CI-enforced
import wall; featurizer
46         single-sourced in `backend/features/rally_seq.py` (train==serve); serving consumes
versioned ARTIFACT BUNDLES only;
47         ship-gate product-side; **pre-committed repo split** (`sports-temporal-models`) at
Gen-2/second-consumer.
48         "Split driven by productionization, not isolation" (owner's framing).
49     3. **M0 executed:** canonical GT loader (`backend/eval/gt_loader.py` ? closed the two-
loader format fork, round-trip
50         tested vs collector export) . **n=6 owned corpus IN the collector SoR** (262 rallies,
Proprietary) . **Gen-0
51         harness** `python -m training.gen0.harness` (train?LOVO?gate?artifact; v0.1.0 built:
**LOVO 0.626**, floor 0.480
52         passed, sign test 5/6 p=0.109 ? ship=False with honest blockers).
53     4. **Gemini battery DONE (the moat is quantified):** clean-ish Mahadevpura ? heuristic
0.657 / Gen-0 0.744 /
54         two-stage refine 0.83 / **wrapper 0.93**; multi-court Kushagra ? heuristic 0.157 /
Gen-0 0.561 / **wrapper 0.66**
55         (FPs = background-game pickups + dead-time merges, the contrast-metric confound).
**Ship target = beat 0.66 on
56         multi-court**; clean footage is commoditized (serve via Gemini). Baselines committed
in `eval_baselines/`.
57         Demo reel exists: `output/MahadevpuraSingles_highlights.mp4`.
58     5. **Gemini ops (hard-won):** prepaid/cold projects burst-throttle with a MISLEADING
"credits depleted" 429 ? ramp
59         gradually; serial works (`ai_max_workers=1` shipped); chunks re-encode to 1280
(`ai_chunk_scale_width`, PR #80);
60         generate retries = 6 attempts exponential+jitter (PR #83); model = `gemini-flash-
latest` alias (2.5-flash is
61         legacy). 503 capacity waves are real ? provider-fallback registry is load-bearing.
62
63     ## Owned-model architecture (for clarity ? owner asked)
64     **Gen-0/Gen-1 = WASB-substrate + OWNED temporal head.** WASB (frozen, MIT code, C3
checkpoint provenance = ship
65     blocker) produces trajectories; the owned IP = featurizer + head + corpus +
harness/gate. **Gen-2 (Qwen3-VL via
66     ms-swift) = pixels, drops WASB** ? that's the heavyweight step, only entered if the
lightweight head plateaus below
67     target. Gen-0 trains in MINUTES on the 2070S ($0). Multisport: WASB transfers to TT
(0.909) but not pickleball
68     (0.308) ? per-sport detectors or Gen-2.
69
70     ## ?? Next steps / open owner decisions (mirror of docs/NEXT_STEPS.md ? read that for
detail)
71     1. **Owner mission in progress: grow the golden corpus** (multi-court badminton first;
?3 matches/sport, TT next;
72         adults-only for training; per video: label ? `import-local --normalize` ? re-run
harness). Sign test reaches
73         significance around n=10 with 9 wins.
74     2. **Set the ship-gate target** (suggested: LOVO F1@tIoU0.5 ? 0.70 on multi-court folds
+ paired-sign p<0.05).
75     3. **M0.4 increment:** ActionFormer/ASFormer into the harness (still lightweight/local)
+ per-video threshold calibration.
76     4. C13 interviews + camera pilot (PMF) ? **field kit SHIPPED (#104):** flyer/visit-
playbook/answer-sheet/pre-registered
77         signals (associate-pay removed per owner; he handles engagement terms); repo rename;
rotate the chat-pasted Gemini
78         key; collector=SoR confirm.
79     5. **UI-alpha track (owner-gated):** PR-1 async jobs (#87) + PR-2 upload/download (#99)
DONE ? **next = PR-3 identity**
80         (`owner_email` scoping, JWT/header auth w/ `DEV_USER_EMAIL` dev fallback, CORS
narrow, rate limit, quota). The
81         PR2_HANDOFF "tomorrow-morning CUJ" runbook is unblocked for the owner.
82     6. **Open owner decisions:** ratify/merge **#106** video-locality model (it reframes

```

```

hosting ? timestamps are the
83     product); whether the in-flight WSL cache-hygiene slice gets PR'd by its owner or
folded into the next run.
84
85     ## Ramp-up kit
86     - [[current-state]] FIRST (dense checkpoints) · `docs/NEXT_STEPS.md` (live plan) ·
`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`
87     (roadmap) · ADR-008 in `docs/archives/decisions/DECISIONS.md` (topology) ·
`docs/CODE_MAP.md` (note its 2026-06-10
88     sweep banner ? several documented behaviors changed).
89     - Code anchors: `training/gen0/harness.py` · `backend/features/rally_seq.py` ·
`backend/eval/gt_loader.py` ·
90     `eval_baselines/*.json` · collector `src/cli/curate.py` (reclassify-licenses, import-
local --normalize/--force).
91     - Data: labels `Annotation Setup\Collect\Golden Labelled\` · trajectories+Roora
`...\Collect\Trajectories\` ·
92     eval manifest `output/human_lovo_manifest.json` (gitignored) · corpus DB = collector
`data/sports_metadata.db`
93     (backup: `data/sports_metadata.backup-2026-06-10.db`, gitignored).
94
95     ## Gotchas (this session's additions)
96     - Merge discipline: gate on `gh pr view --json statusCheckRollup` == SUCCESS (a piped
`gh pr checks --watch` exit code
97     once let a red PR merge ? #75/#76 incident). Collector DB: stash-and-pop, never
`checkout --` (recovered once from
98     a dangling stash). Other sessions may write [[current-state]] concurrently ? insert
checkpoints, never overwrite.
99     - PowerShell `ConvertTo-Json` reformats config.json wholesale ? edit surgically.
`Select-Object -Last N` buffers
100     background output until exit (no interim tail) ? and loses EVERYTHING if the process
is killed (idle/suspend kills
101     background runs silently); pipe long background runs through `Out-File` instead.
102     - **Sibling sessions share this checkout and flip branches mid-session** (the node24
chip did) ? re-verify
103     `git branch --show-current` before EVERY commit; the post-merge "checkout master &&
pull" scripting once left the
104     tree on master and a commit went straight to the default branch (the AI-7 process slip
? don't repeat it).
105     - Windows ffmpeg here lacks NVENC ("No capable devices found") ? use libx264 for host-
side re-encodes.
106     - The dev container claim in CLAUDE.md is stale for THIS machine: cv2/torch/numpy all
present; GPU via WSL.
107

```

#### 499. assistant (2026-06-14T06:02:20.640Z)

The handoff is stale (still pinned to the #127 fusion-measurement track). Let me refresh its top + next-steps so the next session resumes cleanly at execution-policy \*\*P3/P4\*\*.

#### 500. user (2026-06-14T06:02:39.646Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\session-handoff.md has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 501. assistant (2026-06-14T06:02:48.306Z)

Now add the explicit P3/P4 pickup as the top Next-Step + the new code anchors to the ramp-up kit.

#### 502. user (2026-06-14T06:02:56.655Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\session-handoff.md has been updated successfully. (file state is current in your

context ? no need to Read it back)

### 503. assistant (2026-06-14T06:03:04.392Z)

Adding this session's code/doc anchors to the ramp-up kit:

### 504. user (2026-06-14T06:03:11.440Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\session-handoff.md has been updated successfully. (file state is current in your context ? no need to Read it back)

### 505. assistant (2026-06-14T06:03:22.371Z)

The handoff is refreshed ? immediate pickup (execution-policy P3 ? P4) at the top, this session's landings summarized, a concrete P3/P4 resume plan as Next-Step #0, and the code/doc anchors. Let me do a final cleanup check so I leave the repo tidy.

### 506. user (2026-06-14T06:03:26.944Z)

```
=== worktrees (should be only the main checkout) ===
C:/Users/avidu/Projects/badminton-highlight-indexer d7b1f51 [master]
=== current branch + open PRs ===
master
=== this session's merged PRs (mine) ===
ec2002d feat(infra): ExecutionPolicy P1 ? bound the Gemini hang (op_timeout + bounded poll)
(#133)
2673eee docs: execution-policy framework ? per-job, self-scheduling resilience (design) (#132)
c8bdfdd docs: promote Gate-A (measured win) + per-video output layout plan (#131)
2e5892e feat(eval): fix the learned variant ? train-fold threshold tuning + calibration (C2)
(#130)
f680b8d feat(eval): wire C1+C4 honest reporting into experiment.compare() (additive) (#129)
8398e52 docs(archive): first end-to-end A/B experiment on the golden corpus (n=6) (#128)
d7b1f51 feat(eval): honest ?F1 ? bootstrap CI + paired sign test on the fusion lift (#127)
2c8b1ad feat(eval): measurement-honesty helpers (course corrections C1-C4, additive/default-OFF)
(#124)
23672f2 docs: extend analyzers/reconcilers into a multi-scope signal-fusion framework (warm-up +
shuttle-state) (#119)
```

### 507. assistant (2026-06-14T06:03:39.688Z)

All clean: **only the main checkout** (no stray worktrees from this session's nine isolated branches), on `master`, and the only open PR is a sibling's. The handoff and `current-state.md` are both refreshed with the explicit **P3 ? P4** pickup.

It's been a genuinely great session. A quick ledger of what landed, for the record:

```
**Rally-detection quality track**
- A multi-scope **signal-fusion framework** doc + a verified multi-agent **prior-art review**
(#119) ? honest verdict: skeleton is standard, cite don't claim.
- **Measurement-honesty helpers** C1?C4 (#124) + wired into `compare()` (#129).
- The **first golden A/B experiment** (archived, #128) ? **Gate-A promoted** (+19% F1, 6/6
matches, p=0.031, #131) and the **learned variant fixed** (+38% F1, un-zeroed, #130).
- A **fresh GX010125 reel** in its own isolated dir + a per-video-layout plan (#131).

**Resilience track (your new ask)**
- **`docs/EXECUTION_POLICY.md`** (#132) ? your per-job, self-scheduling, no-master architecture,
validated.
- **P1 (#133):** the actual Gemini-hang fix, `run_bounded`/`poll_until`, 100% module coverage,
polls demoted to DEBUG.
- **P2 (#135):** the self-scheduling substrate (`claim_due_job`, `set_job_waiting`).
```

**\*\*Cleanly handed off:\*\*** P3 (executor self-scheduling loop + resumable handoff) and P4 (WSL/ffmpeg) are spec'd as the next two distinct, high-coverage PRs ? wiring, not invention,

since the substrate's already in.

Thank you for the trust, the sharp ideas (shuttle-state, warm-up, the policy design), and for caring about doing it \*honestly\* and \*well\* over doing it fast ? that's what made this fun. Catch you in the next session. ?

**508. user (2026-06-14T06:11:01.066Z)**

can you tell me a good prompt which would make the new session to start off from the next item P3 etc.? Also are there any untracked files in the local directory? I see the UI asking me to mak a PR. I want to keep this directory clean as well

**509. assistant (2026-06-14T06:11:22.652Z)**

Let me check the working tree state ? the "make a PR" nudge is likely because local `master` is behind origin (all my self-merges landed on origin via `gh`, not pulled locally) and/or untracked scratch files.

**510. user (2026-06-14T06:11:27.372Z)**

=== local vs origin/master ===

**master...origin/master [behind 8]**

```
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
=== untracked + modified (excluding gitignored) ===
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/ab_eval_demo.py
?? scratch/ab_n6_lovo.py
?? scratch/arial.ttf
?? scratch/ci_commit_msg.txt
?? scratch/debug_wsl.py
?? scratch/extract_until_done.ps1
?? scratch/extract_wrapper.log
?? scratch/keep_awake.ps1
?? scratch/obs_commit_msg.txt
?? scratch/p3_commit_msg.txt
?? scratch/p3_smoke.py
?? scratch/p4_commit_msg.txt
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/run_gx010125_full.py
?? scratch/run_gx010125_rescore.py
?? scratch/sf_commit_msg.txt
?? scratch/speedup_commit_msg.txt
?? scratch/stitch_min5.py
?? scratch/trim_video.py
=== is output/ and scratch/ gitignored? ===
output/GX010125_run
artifacts/
=== untracked NON-ignored files (what the UI flags) ===
?? artifacts/rally_tal_gen0/0.1.0/manifest.json
?? artifacts/rally_tal_gen0/0.1.0/weights.json
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
```



```
?? scratch/PR2_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
?? scratch/ab_eval_demo.py
?? scratch/ab_n6_lovo.py
?? scratch/arial.ttf
?? scratch/ci_commit_msg.txt
?? scratch/debug_wsl.py
?? scratch/extract_until_done.ps1
?? scratch/extract_wrapper.log
?? scratch/keep_awake.ps1
?? scratch/obs_commit_msg.txt
?? scratch/p3_commit_msg.txt
?? scratch/p3_smoke.py
?? scratch/p4_commit_msg.txt
?? scratch/probe_cv2_count.py
?? scratch/probe_gemini.py
?? scratch/repro_drawtext_font.py
?? scratch/repro_movie_overlay.py
?? scratch/run_gx010125.py
?? scratch/run_gx010125_full.py
?? scratch/run_gx010125_rescore.py
?? scratch/sf_commit_msg.txt
?? scratch/speedup_commit_msg.txt
?? scratch/stitch_min5.py
?? scratch/trim_video.py
```

#### 511. assistant (2026-06-14T06:12:18.873Z)

Two things going on ? let me sync first, then explain the untracked files.

#### 512. user (2026-06-14T06:12:23.846Z)

```
=== sync local master to origin (clears the '8 behind') ===
  create mode 100644 tests/test_job_self_scheduling.py
  create mode 100644 tests/test_run_telemetry.py
  create mode 100644 tests/test_variant_tuning.py
=== is scratch/ in .gitignore? ===
Binary file .gitignore matches
=== which untracked files are NEW this session (mine) ===
scratch/ab_eval_demo.py scratch/ab_n6_lovo.py scratch/run_gx010125_full.py
scratch/run_gx010125_rescore.py
```

#### 513. assistant (2026-06-14T06:13:27.467Z)

Synced ? `git pull --ff-only` pulled the 8 commits, so **local `master` now matches origin**. That "make a PR" nudge was your IDE flagging that local was behind (all my self-merges landed on origin via `gh`, not pulled locally); it should clear now.

### Resume prompt for the new session

Paste this ? it's self-contained (a fresh session needs no memory of this one):

```
> Resume the badminton-highlight-indexer project. First `git pull --ff-only` on master, then read `memory/session-handoff.md` and the top checkpoint of `memory/current-state.md` to orient.
>
> Pick up the EXECUTION_POLICY track at P3 (design: `docs/EXECUTION_POLICY.md`; P1 `backend/infrastructure/execution_policy.py` and P2 `database.py::claim_due_job`/`set_job_waiting` are already merged).
> - P3 (one PR): wire `backend/main.py`'s job executor to a `db.claim_due_job()` self-scheduling loop that honors `WAIT_AND_RESUME` ? when a long op exhausts its policy budget, call `db.set_job_waiting(job_id, delay_sec)` to release + reschedule instead of blocking ? and make the AI handoff resumable (persist per-window confirmation so a resumed handoff only does the unconfirmed windows; this makes the GX010125 reel exhaustive across resumes with no re-spent Gemini calls).
```

```
> - **P4 (next PR):** wrap the WSL/WASB GPU pass + the ffmpeg calls in
`run_bounded`/`poll_until`.
>
> Constraints: one distinct, **self-merged-on-green** PR per phase; **very high coverage**
(fake-clock / fake-hang tests like `tests/test_execution_policy.py`); work in an **isolated git
worktree off `origin/master`** (shared checkout ? branch from `origin/master`, re-verify HEAD
before commit, stage explicit paths).
```

## Untracked files ? yes, and here's the lay of the land

After the pull, \*\*no tracked files are modified\*\* ? the working tree is clean except untracked:

- \*\*`output/...`\*\* (incl. the new `output/GX010125\_run/` + the reel) ? \*\*gitignored\*\* ? won't ever pollute git.
- \*\*`scratch/...`\*\* ? \*\*not\*\* gitignored. Mostly long-standing carryovers (PR handoffs, repro scripts, `commit\_msg` files) the owner has intentionally kept untracked across many sessions; \*\*4 are mine this session\*\* (`ab\_eval\_demo.py`, `ab\_n6\_lovo.py`, `run\_gx010125\_full.py`, `run\_gx010125\_rescore.py` ? the repros referenced in the handoff).
- \*\*`artifacts/rally\_tal\_gen0/0.1.0/`\*\* ? the Gen-0 model bundle (ADR-008), an intentional carryover.

None of these are staged or at risk of an accidental commit (the "stage explicit paths" discipline prevents that). For a \*\*pristine `git status`\*\*, the clean fix is a one-line `.gitignore` PR adding `scratch/` and `artifacts/rally\_tal\_gen0/` ? but those have deliberately been left untracked-not-ignored for many sessions, so I didn't want to change that pattern unilaterally at wrap-up. Want me to:

1. \*\*Add `scratch/` + the gen0 bundle to `.gitignore`\*\* (tiny self-merged PR ? clean status, repros stay on disk), or
2. \*\*Delete just my 4 session repro scripts\*\* (minimal footprint, leaves the rest as-is), or
3. \*\*Leave it\*\* (it's already commit-safe; the nudge was only the sync, now fixed)?

Say the word and I'll do it before you go ? otherwise you're all set. ?

### 514. user (2026-06-14T06:14:55.474Z)

please do so and also I want to cleanup a lot of .md files from the repo root into the docs/ appropriatley. can you suggest that sad well?

### 515. assistant (2026-06-14T06:15:10.312Z)

On it. Let me first see what `.md` files are sitting at the repo root so I can propose a clean mapping, then do the `.gitignore` cleanup.

### 516. user (2026-06-14T06:15:13.111Z)

=== root-level .md files (tracked) ===

CLAUDE.md  
README.md

=== root .md count ===

2

=== current .gitignore tail (to append cleanly) ===

htmlcov/  
.pytest\_cache/

## Local Claude Code config (machine-specific)

.claude/  
.coverage

### 517. assistant (2026-06-14T06:15:28.823Z)

Only \*\*2 tracked root `.md` files\*\* ? `CLAUDE.md` and `README.md` ? and both legitimately belong at the root. So the clutter you're seeing is probably \*\*untracked\*\* `.md` files. Let me check the actual filesystem (root + scratch) to be sure I propose the right thing.

#### 518. user (2026-06-14T06:15:32.703Z)

=== ALL .md at repo root (tracked + untracked) ===

CLAUDE.md  
README.md

=== .md files in scratch/ (untracked clutter) ===

scratch/HIGH\_RISK\_REMEDIATION\_PLAN.md  
scratch/PR1\_BODY.md  
scratch/PR1\_HANDOFF.md  
scratch/PR2\_HANDOFF.md  
scratch/PR88\_FIX.md

=== any other stray .md outside docs/ + backend/ + archives? ===

training/README.md

=== how many .md total tracked, and how many already in docs/ ===

tracked total: 58 | in docs/: 54

#### 519. assistant (2026-06-14T06:16:24.493Z)

Clear picture now: there's actually **nothing** to move from root into `docs/` ? the root has only `CLAUDE.md` + `README.md` (both belong there), `training/README.md` is correctly placed, and `docs/` already holds all 54 doc files. The `.md` "clutter" you're seeing is the **5** untracked `scratch/` notes\* (`PR1\_HANDOFF.md`, `PR2\_HANDOFF.md`, `PR88\_FIX.md`, `HIGH\_RISK\_REMEDIATION\_PLAN.md`, `PR1\_BODY.md`) ? throwaway PR-body/handoff scratch, not docs. The right move is to gitignore `scratch/` (they vanish from status), not relocate them.

Let me verify the `.gitignore` (grep flagged it "Binary" ? worth checking it isn't corrupted) before I append.

#### 520. user (2026-06-14T06:16:27.602Z)

=== .gitignore first bytes (BOM/NUL check) ===

```
00000000 #      B  y  t  e  -  c  o  m  p  i  l  e  d
00000020 /      o  p  t  i  m  i  z  e  d      /      D  L
00000040 L      f  i  l  e  s  \r
```

=== does it already mention scratch / artifacts? ===

(neither present)

=== why does check-ignore say artifacts/ ignored but status shows it? ===

#### 521. assistant (2026-06-14T06:17:11.116Z)

The `.gitignore` is fine (normal CRLF text; grep's "Binary" was a stray non-UTF8 byte in a comment). Neither `scratch` nor `artifacts` is ignored. So the clean fix is to add `scratch/` (clears all the `.md` handoff clutter + my repos from `git status`, keeps the files on disk). Doing it as a tiny PR off `origin/master`.

#### 522. user (2026-06-14T06:17:18.076Z)

HEAD is now at 55df82f feat(obs): run-telemetry black box ? fps + GPU temp/clocks/power + CPU/RAM (rotating ~10KB JSONL) (#134)

=== tail of updated .gitignore ===

■

**Throwaway repro scripts + PR-body / handoff scratch notes (already excluded from ruff + pytest; never meant to be committed ? keep them out of `git status` too).**

scratch/

=== confirm scratch now ignored, tracked scratch files unaffected (should be 0) ===

scratch/ now ignored ?

tracked files under scratch/ (must be 0): 7

#### 523. assistant (2026-06-14T06:18:01.091Z)

There are `**7 *tracked*` files in ``scratch/`**` (committed tools like ``trim_video.py``) ?  
gitignoring ``scratch/`` leaves those tracked (git keeps them) while hiding the untracked clutter,  
which is exactly what we want. Let me commit + self-merge this tiny change.

#### 524. user (2026-06-14T06:18:09.030Z)

```
=== the 7 tracked scratch files (stay tracked) ===
scratch/analyze_motion.py
scratch/clean_state.py
scratch/diagnose_rally.py
scratch/gemini_video_debugger.py
scratch/list_models.py
scratch/test_ai_segmenter.py
scratch/test_thresholds.py
On branch chore/gitignore-scratch
Your branch is ahead of 'origin/master' by 1 commit.
  (use "git push" to publish your local commits)

nothing to commit, working tree clean
* [new branch]      chore/gitignore-scratch -> chore/gitignore-scratch
5fad927 @placeholder@
```

#### 525. assistant (2026-06-14T06:18:15.706Z)

The placeholder slipped in again (shell-quoting). Amending the message properly + force-pushing to my branch.

#### 526. user (2026-06-14T06:18:22.448Z)

```
+ 5fad927...d4e66c7 chore/gitignore-scratch -> chore/gitignore-scratch (forced update)
d4e66c7 chore: gitignore scratch/ to keep git status clean
```

#### 527. user (2026-06-14T06:18:38.806Z)

<https://github.com/avidullu/badminton-highlight-indexer/pull/136>

#### 528. user (2026-06-14T06:20:29.848Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    18s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490471073/job/81254427802
Full test suite (offline, mocked)                pass    1m30s  https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490471073/job/81254427801
Mypy (lenient baseline) pass    31s  https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490471073/job/81254427829
Ruff lint      pass    14s  https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490471073/job/81254427814
Syntax sweep + import smoke (no GPU stack)      pass    18s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490471073/job/81254427802
```

#### 529. user (2026-06-14T06:20:41.120Z)

```
MERGED
wt removed
Deleted branch chore/gitignore-scratch (was d4e66c7).
  1 file changed, 0 insertions(+), 0 deletions(-)
=== resulting git status (should be clean but for artifacts/) ===
?? artifacts/
```

#### 530. assistant (2026-06-14T06:21:16.702Z)

Done ? and the result is clean. After merging [#136](<https://github.com/avidullu/badminton-highlight-indexer/pull/136>) and pulling, ``git status`` now shows `**only `?? artifacts/`**` ? all the ``scratch/`` clutter (including those 5 stale ``PR*_HANDOFF.md`` / ``HIGH_RISK_*.md`` notes + my repros) is now hidden, while the 7 tracked scratch tools stay tracked.

## On the `.md` cleanup ? honest finding

There were actually **no** loose `.md` files at the repo root to move into `docs/`:

- **Root:** only `CLAUDE.md` + `README.md` ? both **belong** at the root (project instructions + readme). Leave them.
- **docs/** already holds all 54 doc files, well-organized; `training/README.md` is correctly placed in its subdir.
- The `.md` clutter you were seeing = the **scratch/` handoff/PR-body notes**, which aren't docs ? they're throwaway. The `scratch/` gitignore (just merged) makes them disappear from `git status`. ?

So there's nothing to relocate root?`docs/`; the cleanup you wanted is the scratch one, now done.

## One item left: `artifacts/`

The last untracked thing is **artifacts/rally\_tal\_gen0/0.1.0/`** (`manifest.json` + `weights.json`) ? the **Gen-0** model bundle (ADR-008). This one's a genuine owner decision, because ADR-008 says the **manifest** should be **committed** (provenance) and only the **weights** **gitignored**. So I didn't touch it. Two ways to make status fully pristine ? your call (say the word and I'll do it as a tiny PR):

1. **Follow ADR-008:** commit `manifest.json`, gitignore `weights.json` (keeps the provenance record, hides the binary).
2. **Gitignore the whole** `artifacts/rally\_tal\_gen0/` (simplest pristine status; defers the ADR-008 manifest-commit).

Otherwise ? local `master` is synced, working tree is clean but for that one bundle, the handoff + resume prompt are ready, and you're all set to pick up at **execution-policy P3** next session. Genuinely a pleasure. ?

### 531. user (2026-06-14T06:21:19.110Z)

these run eval python files? they can be segregated away?

### 532. assistant (2026-06-14T06:21:33.094Z)

Let me see exactly what's there ? the root-level `run\_\*.py` entry scripts and the `backend/eval/` drivers ? so I can propose a clean segregation.

### 533. user (2026-06-14T06:21:36.595Z)

```
=== root-level .py files ===
run_all_tests.py
run_eval.py
run_phase3_eval.py
run_process_and_stitch.py
run_regression.py
setup.py
```

```
=== backend/eval/ contents (drivers vs library) ===
__init__.py
annotations.py
audio/__init__.py
audio/e1_probe.py
batch.py
calibrate_local.py
calibrate_wasb.py
calibration.py
classifier.py
distill_local.py
eval_partial_labels.py
eval_stats.py
experiment.py
export_wasb_segments.py
fusion_audio.py
fusion_compare.py
```

```
fusion_golden.py
gemini_refine.py
gt_loader.py
harness.py
metrics.py
rally_seq_proto.py
regression.py
score_calibration.py
segmentation_metrics.py
training.py
```

=== which root/eval .py have a \_\_main\_\_ / argparse (i.e. runnable CLI drivers)? ===

```
backend/eval/audio/e1_probe.py
backend/eval/calibrate_local.py
backend/eval/calibrate_wasb.py
backend/eval/distill_local.py
backend/eval/eval_partial_labels.py
backend/eval/export_wasb_segments.py
backend/eval/fusion_audio.py
backend/eval/fusion_compare.py
backend/eval/fusion_golden.py
backend/eval/gemini_refine.py
backend/eval/rally_seq_proto.py
backend/main.py
backend/pipeline/detectors/wasb_infer.py
backend/tools/rally_slicer.py
run_all_tests.py
run_eval.py
run_phase3_eval.py
run_process_and_stitch.py
run_regression.py
scratch/analyze_motion.py
scratch/clean_state.py
scratch/diagnose_rally.py
scratch/gemini_video_debugger.py
scratch/test_ai_segmenter.py
scratch/test_thresholds.py
setup.py
tools/cleanup_caches.py
training/gen0/harness.py
```

#### 534. assistant (2026-06-14T06:22:23.371Z)

Good question ? and the honest answer is "some yes, most no," because of imports. Let me confirm which are safe to move vs. which are library code in disguise.

#### 535. user (2026-06-14T06:22:28.832Z)

=== are the 3 root run-drivers imported anywhere (or pure scripts)? ===

```
run_eval.py: imported by 1 file(s)
run_phase3_eval.py: imported by 0 file(s)
run_process_and_stitch.py: imported by 0 file(s)
```

=== ruff.toml / CI references to the root run scripts ===

```
ruff.toml:20:"run_phase3_eval.py" = ["E402"]
ruff.toml:21:"run_process_and_stitch.py" = ["E402"]
.github/workflows/ci.yml:85:      # pytest.importorskip ? run_all_tests.py prints which. The
coverage floor
.github/workflows/ci.yml:92:      run: python run_all_tests.py --cov=backend --cov=training
--cov-report=term --cov-fail-under=70
```

=== which backend/eval drivers are IMPORTED by non-eval code (can't move) ===

```
distill_local: imported by 1 non-test/non-self file(s)
calibrate_wasb: imported by 8 non-test/non-self file(s)
gemini_refine: imported by 2 non-test/non-self file(s)
fusion_compare: imported by 1 non-test/non-self file(s)
fusion_golden: imported by 2 non-test/non-self file(s)
```

calibrate\_local: imported by 0 non-test/non-self file(s)

### 536. assistant (2026-06-14T06:23:05.783Z)

Here's the honest read ? and it splits cleanly into "yes, safe" and "no, don't."

#### **`backend/eval/\*.py` ? leave them (they're library, not scripts)**

Despite having `\_\_main\_\_`/argparse, almost all are **imported across the codebase**, so they're library+CLI hybrids, not loose scripts:

- `calibrate\_wasb` ? imported by **\*\*8\*\*** files (its `load\_golden\_gts` is the canonical golden reader)
- `gemini\_refine` (2, incl. `experiment.py`), `fusion\_golden` (2), `distill\_local` (1), `fusion\_compare` (1)?

Moving these would break the import graph for no real gain ? they're correctly cohesive as the `eval` package. **Recommend: don't touch `backend/eval/`.** (Only `calibrate\_local` has 0 importers, not worth moving one file.)

#### **Root `\*.py` ? 3 stay, 3 can move**

Only 6 `.py` at root. Three are **canonical entry points** referenced elsewhere ? keep at root:

- `setup.py` (packaging) · `run\_all\_tests.py` (**\*\*CI runs it\*\***, `ci.yml:92`) · `run\_regression.py` (the documented promotion-gate CLI ? moving needs doc edits)

Three are **run-drivers**, safe to segregate into `scripts/` (low/no coupling):

file	importers	move cost
---	---	---
`run_phase3_eval.py`	<b>**0**</b>	+ update 1 `ruff.toml` E402 path
`run_process_and_stitch.py`	<b>**0**</b>	+ update 1 `ruff.toml` E402 path
`run_eval.py`	1	+ fix that 1 import

So: `git mv` those 3 into `scripts/`, update the 2 `ruff.toml` per-file-ignore paths + the 1 import, done. Root drops from 6 ? 3 `.py`, all of which genuinely belong there.

**My recommendation:** do the 3-file `scripts/` move (tidy, low-risk), leave `backend/eval/` and the 3 canonical entry points alone. Want me to spin that up as a small self-merged PR? (If you'd also like `run\_regression.py` moved, I'll do it too ? just costs a couple of doc-reference updates in `EXPERIMENT\_HARNESS.md`/`NEXT\_STEPS.md`.)

### 537. user (2026-06-14T06:23:13.572Z)

additionally can you add to memory and docs as to where to add new scripts from the future. I believe the directory layout is written in the repo somewhere but it may be stale or could be heavily improved with guidelines which you and others can read before the committing files anywhere

### 538. assistant (2026-06-14T06:23:52.106Z)

On it. Let me check the `run\_eval.py` importer (to know the move cost) and whether `CODE\_MAP.md` already has a layout section to refresh.

### 539. user (2026-06-14T06:23:55.288Z)

=== who imports run\_eval? ===

```
tests/test_eval_harness.py:11:import run_eval
```

=== does a scripts/ dir already exist? ===

```
ls: cannot access 'scripts/': No such file or directory
tools/
```

=== CODE\_MAP.md: any directory-layout / 'where to put' section? ===

```
1:# CODE_MAP.md
```

5:**Regenerate:** rerun the `build-code-map` workflow (parallel subsystem readers ? this doc).  
Re-run after any registry/contract/WSL-interface change, or after a directory restructure (this

```

revision follows #40).
20:## 1. PIPELINE
22:### 1A. Runtime: video file -> highlight reel
50:### 1B. WASB shuttle substrate (P2 of wasb_hybrid)
62:### 1C. Calibration / eval flow (NO pipeline run unless told)
80:## 2. SUBSYSTEMS
82:### 2.0 Orchestration & config
108:### 2.1 Segmenters (under pipeline/segmenters) ?
117:- `@backend/pipeline/segmenters/yolo_hybrid.py::HybridYoloSegmenter` (name `yolo_hybrid`,
no YOLO remains; ADR-005). **Stage-1 detection EXTRACTED to `detectors/person_detector.py`
(fusion P1, no behavior change)** ? the segmenter now ORCHESTRATES only (chunking, pipelined-vs-
sequential P2 prefetch where ffmpeg overlaps sequential GPU, AI-provider handoff, candidate
persist/link) and delegates detection to `self.person` (`PersonDetector`). Imports
`_DEFAULT_DETECTOR` from person_detector for the `detection_params` snapshot. ?
`yolo_tracker`/`bytetrack.yaml` is dead/back-compat; `gemini_*` config keys are deprecated
aliases still honored; still mypy-quarantined for pre-existing orchestration debt (the
torchvision code that caused it left).
121:### 2.2 Detectors & windowing (WSL-quarantined; ABC abstraction complete)
128:- `@backend/pipeline/detectors/wasb_infer.py` (WSL) ? `::run` core. ? imports WASB-SBDT repo
modules + torch ? NOT importable on Windows. Reboot-resumable detector cache:
`::DET_FILE='det_raw.jsonl'`, `::MANIFEST_FILE='manifest.json'`, `::CACHE_VERSION=1`,
`::load_and_clean_cache` (reads longest CONTIGUOUS prefix where `w==line index`, rewrites to
heal crash-truncated tail), `::manifest_compatible` (? `frames_dir` abspath'd but `weights`
compared raw), `::atomic_write_text/_atomic_write_trajectory`, `::dets_to_tracker` (xy MUST be
np.array), `::build_cfg` (`runner.gpus=[0]` hardcoded), `::extract_frames` (cv2-PNG SLOW),
`::default_cache_dir` (`<stem>_wasbcache` next to out). ? GPU pass resumable; CPU tracker
STATEFUL -> always fully re-run; `--fresh` ignores cache.
133:### 2.3 Eval & calibration (all pure core; I/O in CLI drivers)
152:### 2.3a Audio (E1 probe; NOT adopted for fusion ? ADR-007)
156:### 2.4 Stitcher / tools / utils
158:- `@backend/tools/rally_slicer.py` ? standalone CLI. `::compute_clip_windows` (pure, unit-
tested), `::load_rally_labels` (golden schema; ? silently skips degenerate rows),
`::slice_rallies`, `::ClipWindow` (rally:str), `::BEGIN_PADDING=0.5`/`::END_PADDING=1.0` (?
duplicate compiler's by copy). `read_time` -> backend/eval/annotations.parse_timestamp`.
159:- `@tools/cleanup_caches.py` ? **top-level** `tools/` ops tool (NOT backend runtime). Dry-
run-first cache GC. Pure `::classify` (name,is_dir,location)/`::plan_deletions` (unit-tested =
the destructive decision path) + I/O `::scan_wsl` (du+find; ? arg-globbing NOT a `for` loop ?
the login shell empties loop vars) / `::scan_windows` / `::delete`. CLI: `python -m
tools.cleanup_caches` (report) . `--apply` / `--keep-video <id|stem>` / `--older-than N` /
`--include-wasbcache` / `--summary` (hook nudge). Purges `*_frames`/staged/proxy/handoff/tmp;
keeps CSV/DB/reels. Companion to the runner self-clean (2.2). See `[[cache-hygiene-policy]]`
memory.
163:### 2.5 Validators ?
173:### 2.6 Sports / Providers / Annotations ?
181:### 2.7 Infrastructure (DB) ?
184:### 2.8 Reporting / observability ? (per-video, soft-dep `obsreport`)
187:- `@backend/reporting/spec.yaml` ? $ declarative report layout + footgun engine consumed by
obsreport. `::customer_verdict` (pass/pass_with_warnings/fail messages). `::sections` (ordered
`{id,title,kind,audience?{both,customer,technical}}`). ? `::rules` ? each
`{code,severity,when[{path,op,value}],message,faq_ref,customer_message?}`; `when` clauses AND'd;
`path` dotted-navigates the recorder JSON. ? `path` strings are a HARD coupling to harvest's
stage/metric/detail names ? renaming a detail in harvest.py silently breaks the matching rule.
Rules: `silent_provider_noop` (error: no api key + 0 segs), `oversize_clip_skipped` (error:
provider refused chunks over MAX_UPLOAD_MB ? 4K stream-copy footgun),
`ai_empty_vs_no_rallies` (warn), `fps_fallback_used` (warn), `synthetic_motion_metadata` (warn),
`segmenter_divergence` (info), `reingest_replaced_prior` (info), `validation_failed` (error: halts
indexing).
192:## 3. DATA CONTRACTS (canonical shapes ? one place)
230:## 4. WASB WSL CONTRACT (external dep `~/models/WASB-SBDT`, MIT; NOT in this repo ? claims
below are unverified against external source, only the repo-side caller wasb_infer.py was
checked)
240:## 5. CROSS-REPO (sibling private repo `sports-data-collector` ? external, unverified here)

```



`run\_eval.py` is imported by a test (so it's coupled ? I'll grandfather it at root), but  
`run\_phase3\_eval.py` + `run\_process\_and\_stitch.py` have zero importers, so they move cleanly.  
`scripts/` doesn't exist yet, `tools/` does (ops tools), and `CODE\_MAP.md` has subsystem  
sections but **\*\*no "where do new files go" guideline\*\*** ? exactly the gap you flagged. Let me  
check doc/code references to the 2 scripts before moving.

#### 541. user (2026-06-14T06:24:56.426Z)

```
=== references to the 2 scripts I'll move (docs + code, excl. the files + scratch) ===
backend/eval/batch.py:10:or any API. The orchestration/side-effects live in
`run_phase3_eval.py`.
backend/main.py:503:    # path too ? parity with run_process_and_stitch.py. Segments WITHOUT
shot_count
backend/utils/hashing.py:4:(`backend/main.py`, `run_eval.py`, `run_phase3_eval.py`,
`run_process_and_stitch.py`)
docs/CODE_MAP.md:76:Entrypoints: `@run_eval.py` (single video score), `@run_phase3_eval.py`
(manifest batch, can segment), `@run_regression.py` (baseline gate, exit 1=regression),
`@backend/eval/calibrate_wasb.py` + `@backend/eval/export_wasb_segments.py` (WASB tuning
drivers).
docs/CODE_MAP.md:90:- `@run_process_and_stitch.py` ? one-shot demo; ? hardcoded paths,
interactive `input()` (blocks automation), skips validation. Config via
`backend.config.load_config` (typed); segmenter, padding, and db path all read from the model
(no literal fallbacks).
docs/CODE_MAP.md:92:- `@run_phase3_eval.py::main` ? walk manifest, segment-if-needed, score,
aggregate. `load_dotenv()` at import. ? DB key = file MD5, NOT manifest's `video_id` (kept as
label); config via `backend.config.load_config`; `segmenter_cls(db, cfg)` receives the typed
`AppConfig` (same object the live pipeline feeds segmenters).
docs/CODE_MAP.md:140:- `@backend/eval/batch.py` ? corpus aggregation (pure; orchestration is in
run_phase3_eval). `::aggregate` (micro pool TP/FP/FN, TP-weighted mean_iou, macro F1),
`::default_stages_for` (`auto`; `::FINAL_ONLY_SEGMENTERS={motion,gemini}`),
`::resolve_video_path` (normalizes collector `./data/..` Windows paths), `::format_aggregate`.
docs/CODE_MAP.md:157:- `@backend/pipeline/stitcher/compiler.py::VideoCompiler` ?
`compile_highlights(raw, segments:[(start,end)], out_name)`. Per-clip re-encode -> `-f concat -c
copy`. `::parse_time` (mirrors gemini.parse_time; ? unparseable -> 0.0),
`::get_video_duration`. ? if duration probe fails (0.0) NO boundary validation;
`begin_padding`/`end_padding` ctor defaults 0.5/1.0 but every live call site (`/api/compile`,
CLI trim, `run_process_and_stitch.py`) threads `config.json
stitching.{begin_padding,end_padding}` (typed default 0.5/0.5); temp clips in a
TemporaryDirectory under output_dir; raises ValueError/FileNotFoundError/RuntimeError; uses
`print()`. ? two-stage codec contract: clips re-encoded with IDENTICAL settings precisely so
final concat can stream-copy ? change a preset and concat may break. ? optional **branding
watermark** (`stitching.watermark`, default-on): `::watermark_filter` builds a `-vf` (drawtext
via `textfile=` + optional logo `movie/overlay`) applied to the per-clip re-encode only (concat
stays `-c copy`); a failed watermark encode retries the clip WITHOUT it (never nukes the reel,
and logs the ffmpeg stderr reason). Pass `watermark=` to `VideoCompiler(...)` (a
`WatermarkConfig` model or dict; `None` = off). Default font = bundled Apache-2.0
`::BUNDLED_FONT` (`backend/assets/fonts/RobotoSlab-Regular.ttf`) so it renders without
fontconfig; `font_path` overrides.
docs/CODE_MAP.md:242:- `python -m src.cli.curate export` (collector) ->
`<out>/<sport>/<video_id>.csv` (indexer `start,end` format) + `manifest.json` -> consumed by
`@backend/eval/batch.py` + `@run_phase3_eval.py`.
docs/CODE_MAP.md:282:| Filter highlights by shot count | `config.json stitching.min_shot_count`
? enforced on BOTH `@backend/main.py::compile_highlights` (`/api/compile`) and
`@run_process_and_stitch.py` (segments without `shot_count` metadata are exempt) |
=== how test_eval_harness uses run_eval (to confirm it stays at root) ===
    evaluate, get_predictions, format_report_text, report_to_dict, _rows_to_intervals,
)
import run_eval
```

@pytest.fixture

#### 542. assistant (2026-06-14T06:26:22.021Z)

The references are mostly **\*\*doc/comment mentions\*\*** (CODE\_MAP entrypoints, a few docstrings) plus

the 2 mandatory `ruff.toml` E402 paths ? all updatable. `run\_eval.py` is test-imported so I'll grandfather it at root. Let me execute: a worktree, move the 2 clean scripts, fix the references, and add the **\*\*\*"where new files go" guidelines\*\*** to CODE\_MAP (the durable part of your ask).

#### 543. user (2026-06-14T06:26:51.239Z)

```
fatal: unable to access 'https://github.com/avidullu/badminton-highlight-indexer.git/': Failed
to connect to github.com port 443 after 21118 ms: Couldn't connect to server
HEAD is now at fa12829 chore: gitignore scratch/ to keep git status clean (#136)
=== update references in ruff.toml + CODE_MAP + 3 docstrings (filename -> scripts/ path) ===
--- ruff.toml per-file-ignores now:
20:"scripts/run_phase3_eval.py" = ["E402"]
21:"scripts/run_process_and_stitch.py" = ["E402"]
--- sanity: no accidental double scripts/scripts/:
```

#### 544. assistant (2026-06-14T06:27:02.344Z)

Worktree built off the latest local `origin/master` (the fetch hit a transient network blip ? I'll retry the push later); `git mv` + ruff.toml updates are clean. Let me read CODE\_MAP's top to slot in the guidelines section and verify the sed didn't mangle it.

#### 545. user (2026-06-14T06:27:03.044Z)

```
1      # CODE_MAP.md
2
3      **Purpose:** cold-start navigation map for the badminton-highlight-indexer ? a self-
hosted video ? highlight-reel pipeline (CV/AI rally detection ? SQLite ? ffmpeg stitch) with a
segmenter-agnostic eval/calibration stack, typed config, and per-video observability reporting.
4
5      **Regenerate:** rerun the `build-code-map` workflow (parallel subsystem readers ? this
doc). Re-run after any registry/contract/WSL-interface change, or after a directory restructure
(this revision follows #40).
6
7      ? **2026-06-10 audit-hardening sweep changed several behaviors documented below ? trust
git history over stale lines until the next full regen:** `add_video` is now an **UPSERT** (no
FK-cascade wipe on re-ingest; segmentation is destroy-AFTER-confirm via
`segment_watermarks`/`prune_segments`); detector artifacts are keyed by **video_id**, not
basename; **`indexing.wasb.*` is now a real typed config** (was silently dropped); `timeout_sec`
defaults **1800s** (was 0/hang-forever); SQLite opens with **WAL + busy_timeout**;
`match_intervals` requires real overlap (iou>0); `run_regression` baseline lives in tracked
`eval_baselines/` and **exits 3 on a missing baseline**; the API validates
`video_path`/`output_filename`. New tests: `test_api_endpoints`, `test_compiler`,
`test_wasb_runner`, `test_reingest_safety`, `test_detector_keying`, `test_config_wasb`,
`test_path_safety`, `test_reliability_hardening` (+ collector `test_licensing_gate`,
`test_import_guard`, `test_cvat_robustness`).
8
9      **LEGEND**
10     - `@path` = file (always repo-relative)
11     - `:::sym` = symbol (class/func/const) in the nearest preceding `@path`
12     - `->` = dataflow / call / "produces"
13     - `$` = data contract (canonical shape)
14     - `?` = gotcha / footgun
15     - `?` = registry / plugin extension point
16     - `WSL` = runs inside WSL2 conda env, NOT the Windows Python process
17
18     ---
19
20     ## 1. PIPELINE
```

#### 546. assistant (2026-06-14T06:27:19.085Z)

CODE\_MAP's top is intact. Now I'll add the **\*\*\*\$0 Repo layout & where new files go\*\*\*** guidelines section (the durable part of your ask) right after the legend, plus a `scripts/README.md`.

#### 547. user (2026-06-14T06:27:35.688Z)

The file C:/Users/avidu/Projects/bhi-scripts-wt/docs/CODE\_MAP.md has been updated successfully.  
(file state is current in your context ? no need to Read it back)

#### 548. assistant (2026-06-14T06:27:43.627Z)

Now the `scripts/README.md`, then verify ruff/imports before committing.

#### 549. assistant (2026-06-14T06:27:49.234Z)

### scripts/

**\*\*Standalone run-drivers / one-off entry CLIs\*\*** ? scripts with a `\_\_main\_\_` that you *\*run\**, not *\*import\**. Run as `python scripts/<name>.py`.

What belongs here (vs. elsewhere ? see `docs/CODE\_MAP.md` §0 for the full table):

- ? a CLI that orchestrates a pipeline/eval run end-to-end (e.g. `run\_phase3\_eval.py`, `run\_process\_and\_stitch.py`)
- ? **\*\*importable eval library\*\*** (metrics, calibration, a scorer) ? `backend/eval/`
- ? an **\*\*ops/maintenance tool\*\*** run as `python -m tools.x` ? `tools/`
- ? a **\*\*throwaway repro / debug\*\*** script ? `scratch/` (gitignored)
- ? anything imported by app code ? it's library, put it under `backend/`

If a script needs `load\_dotenv()` / `sys.path` setup before importing `backend`, add  
`scripts/<name>.py" = ["E402"]` to `ruff.toml` (import-order is load-bearing there).

Grandfathered at the repo root (referenced by CI/packaging/tests ? left in place): `setup.py`,  
`run\_all\_tests.py`, `run\_regression.py`, `run\_eval.py`. **\*\*New\*\*** run-drivers go here, not the  
root.

#### 550. user (2026-06-14T06:27:49.246Z)

File created successfully at: C:/Users/avidu/Projects/bhi-scripts-wt/scripts/README.md (file  
state is current in your context ? no need to Read it back)

#### 551. user (2026-06-14T06:27:59.312Z)

```
=== ruff (the moved scripts must still get their E402 ignore at the new path) ===
All checks passed!
=== moved files present + old paths gone ===
scripts/run_phase3_eval.py
scripts/run_process_and_stitch.py
ls: cannot access 'run_process_and_stitch.py': No such file or directory
=== git status (renames + edits) ===
  M backend/eval/batch.py
  M backend/main.py
  M backend/utils/hashing.py
  M docs/CODE_MAP.md
  M ruff.toml
R   run_phase3_eval.py -> scripts/run_phase3_eval.py
R   run_process_and_stitch.py -> scripts/run_process_and_stitch.py
?? scripts/README.md
=== confirm the test that imports run_eval is unaffected (run_eval stayed at root) ===
run_eval.py
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

#### 552. assistant (2026-06-14T06:28:14.025Z)

Ruff passes (moved scripts keep their E402 ignore at the new path), renames detected,  
`run\_eval.py` grandfathered at root, and the backend edits are docstring/comment-only (no  
behavior change). Committing + pushing (retrying past the network blip).

### 553. user (2026-06-14T06:28:23.000Z)

```
warning: in the working copy of 'backend/eval/batch.py', LF will be replaced by CRLF the next
time Git touches it
warning: in the working copy of 'backend/main.py', LF will be replaced by CRLF the next time Git
touches it
warning: in the working copy of 'backend/utils/hashing.py', LF will be replaced by CRLF the next
time Git touches it
warning: in the working copy of 'docs/CODE_MAP.md', LF will be replaced by CRLF the next time
Git touches it
warning: in the working copy of 'ruff.toml', LF will be replaced by CRLF the next time Git
touches it
warning: in the working copy of 'scripts/README.md', LF will be replaced by CRLF the next time
Git touches it
To https://github.com/avidullu/badminton-highlight-indexer.git
* [new branch]      chore/scripts-dir-and-layout-guidelines -> chore/scripts-dir-and-layout-
guidelines
```

### 554. assistant (2026-06-14T06:28:39.700Z)

Push succeeded. The sed rewrote those files as LF ? let me verify that didn't cause a whole-file line-ending churn in the diff (it shouldn't, since `autocrlf` normalizes blobs to LF) before opening the PR.

### 555. user (2026-06-14T06:28:40.051Z)

```
=== diff stat vs master (edits should be SMALL ? not whole-file CRLF churn) ===
 backend/eval/batch.py      |  2 +-
 backend/main.py            |  2 +-
 backend/utils/hashing.py   |  2 +-
 docs/CODE_MAP.md           | 44 ++++++-----
 ruff.toml                  |  4 +-
 scripts/README.md          | 17 ++++++++
 run_phase3_eval.py => scripts/run_phase3_eval.py |  0
 .../run_process_and_stitch.py |  0
 8 files changed, 57 insertions(+), 14 deletions(-)
=== confirm the actual content edits in the 3 code files (should be 1 line each) ===
-or any API. The orchestration/side-effects live in `run_phase3_eval.py`.
+or any API. The orchestration/side-effects live in `scripts/run_phase3_eval.py`.
-  # path too ? parity with run_process_and_stitch.py. Segments WITHOUT shot_count
+  # path too ? parity with scripts/run_process_and_stitch.py. Segments WITHOUT shot_count
-(`backend/main.py`, `run_eval.py`, `run_phase3_eval.py`, `run_process_and_stitch.py`)
+(`backend/main.py`, `run_eval.py`, `scripts/run_phase3_eval.py`,
`scripts/run_process_and_stitch.py`)
```

### 556. assistant (2026-06-14T06:28:47.583Z)

Clean ? the 3 code files changed exactly 1 line each (the reference update), no line-ending churn, and the 2 files are pure renames. Opening the PR, then self-merge on green.

### 557. user (2026-06-14T06:30:52.014Z)

```
https://github.com/avidullu/badminton-highlight-indexer/pull/138
--- watching CI ---
Syntax sweep + import smoke (no GPU stack)      pass    26s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490675459/job/81255022374
Full test suite (offline, mocked)                pass    1m31s  https://github.com/avidullu/badminton-
highlight-indexer/actions/runs/27490675459/job/81255022340
Mypy (lenient baseline) pass    29s  https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27490675459/job/81255022358
Ruff lint      pass    10s  https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27490675459/job/81255022333
Syntax sweep + import smoke (no GPU stack)      pass    26s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490675459/job/81255022374
```

**558. assistant (2026-06-14T06:30:59.401Z)**

All green. Merging #138 + cleanup, then the memory file you asked for.

**559. user (2026-06-14T06:31:07.249Z)**

```
MERGED
wt removed
Deleted branch chore/scripts-dir-and-layout-guidelines (was 60074ae).
Already up to date.
=== root .py now (cleaner) ===
run_all_tests.py
run_eval.py
run_phase3_eval.py
run_process_and_stitch.py
run_regression.py
setup.py
```

**560. assistant (2026-06-14T06:31:20.665Z)**

`git pull` said "up to date" but the root still shows the moved files ? the network blip left the local `origin/master` ref stale, so the pull no-op'd. Let me fetch explicitly and sync.

**561. user (2026-06-14T06:31:25.900Z)**

```
From https://github.com/avidullu/badminton-highlight-indexer
* branch          master      -> FETCH_HEAD
=== origin/master HEAD (should be #138) ===
8836dfe chore: scripts/ for run-drivers + CODE_MAP $0 file-placement guidelines (#138)
=== local vs origin ===
docs/i3-fusion-measured...origin/docs/i3-fusion-measured
Already up to date.
=== root .py after real sync ===
run_all_tests.py
run_eval.py
run_phase3_eval.py
run_process_and_stitch.py
run_regression.py
setup.py
=== scripts/ ===
```

**562. assistant (2026-06-14T06:31:45.977Z)**

Important: the shared checkout got flipped to a **\*\*sibling's branch\*\*** (`docs/i3-fusion-measured`) by a concurrent session ? that's why the local view still shows the old root layout. **\*\*`origin/master` is correctly tidied\*\*** (#138 is merged at `8836dfe`, `scripts/` has the moved files there). I won't switch the branch (it's the sibling's active work). It'll reflect the new layout once the checkout returns to master.

Now the **\*\*memory\*\*** part of your ask ? a placement convention future sessions read before committing files:

**563. assistant (2026-06-14T06:31:58.497Z)**

```
---
name: file-placement-conventions
description: "Where new files go in the repo ? read BEFORE committing a new file anywhere.
Authoritative table is docs/CODE_MAP.md $0."
metadata:
  type: reference
---

**Before adding/committing a new file, decide its home by KIND (don't dump scripts at the repo root).**
```

The authoritative, always-current table is `**`docs/CODE_MAP.md`` §0 "REPO LAYOUT ? where new files go" (added 2026-06-14, PR #138) ? read it. Quick rules:

- **Backend / runtime code** ? ``backend/<subsystem>/``  
(``pipeline/{segmenters,detectors,validators}``, ``providers``, ``sports``, ``annotations``, ``infrastructure``, ``reporting``, ``stitcher``, ``utils``, ``config``, ``features``, ``api``). Most additions = **register a class in the relevant `?` registry**, not a new module tree.
- **Importable eval LIBRARY** (metrics, calibration, a scorer, a feature ? pure, no import-time I/O) ? ``backend/eval/``. These are imported widely (``calibrate_wasb`` by 8 sites) ? they are library, **do NOT** move them to `scripts/`.
- **Standalone run-driver / one-off entry CLI** (has ``__main__``, NOT imported by app code) ? **scripts/** (created PR #138). Run ``python scripts/<name>.py``. If it does ``load_dotenv()``/``sys.path`` before importing backend, add ``scripts/<name>.py" = ["E402"]`` to ``ruff.toml``.
- **Ops / maintenance tool** (not the served app ? cache GC, batch admin) ? ``tools/``, run ``python -m tools.<name>``.
- **Throwaway repro / debug** ? ``scratch/`` (**gitignored** as of PR #136; excluded from `ruff+pytest`).
- **Test** ? ``tests/test_*.py`` (70% coverage floor on the full-suite lane).
- **Doc** ? ``docs/`` (+ index in ``docs/README.md``); only terminal-state work ? ``docs/archives/``.
- **Model artifact** ? ``artifacts/<name>/<ver>/`` (ADR-008: manifest committed, weights gitignored). **Pipeline outputs** (reels/trajectories/DBs) ? ``output/`` (gitignored).
- **Training code** ? ``training/`` (behind the CI import wall, ADR-008).

**Grandfathered at the repo ROOT** (referenced by CI/packaging/tests ? leave them): ``setup.py``, ``run_all_tests.py`` (CI), ``run_regression.py`` (no-regression gate), ``run_eval.py`` (imported by ``tests/test_eval_harness.py``). **New** run-drivers go in ``scripts/``, never the root.

**Why:** keeps the root clean + makes "where does this go?" a 5-second lookup instead of a guess; the root was accumulating loose ``run_*.py``. See `[[code-map-artifact]]` (CODE\_MAP is the cold-start map) and the cleanup PRs #136 (gitignore scratch/) + #138 (scripts/ + §0 guidelines).

#### 564. user (2026-06-14T06:31:58.522Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\file-placement-conventions.md (file state is current in your context ? no need to Read it back)

#### 565. assistant (2026-06-14T06:32:09.834Z)

Now a one-line pointer in ``MEMORY.md`` (the index loaded each session). Let me read it to place the entry.

#### 566. user (2026-06-14T06:32:10.385Z)

```
1      # Memory Index
2
3      ## ? Start here
4      - [Current state](current-state.md) ? live micro-checkpoint (updated every
checkpoint): open PRs, what just landed, next steps, open decisions. Read this first.
5      - [Session handoff](session-handoff.md) ? current state · next steps · open decisions ·
ramp-up kit. AUTO-refreshed at meaningful checkpoints (PR merged, key decision, phase shift,
session wrap) ? see working-agreement-checkpoints staleness tripwire.
6      - [Code map artifact](code-map-artifact.md) ? for CODE/architecture ramp-up, read
`docs/CODE_MAP.md` FIRST (dense anchor-rich map; replaces a file crawl). How to regenerate it.
7
8      - [Working agreement: attribution](working-agreement-attribution.md) ? never assert
fabricated facts; NEVER invent a provenance (a discussion/review/source that didn't happen). The
silver-bullet rule.
9      - [Working agreement: checkpoints](working-agreement-checkpoints.md) ? user prefers
frequent small PRs + a Claude-consumable current-state file updated at every micro-checkpoint.
10     - [Feedback: tail long-job logs](feedback-tail-long-job-logs.md) ? proactively tail
progress/summaries from long-running background jobs while they run, not just the final result.
11     - [Working agreement: merging](working-agreement-merging.md) ? owner authorized Claude
```

to MERGE its own good-engineering-fix PRs on verified-green CI (2026-06-12);  
product/strategy/licensing stay owner-gated.

12 - [Gotcha: shared-checkout branch collisions](gotcha-shared-checkout-branch-collisions.md) ? spawned tasks + sibling sessions SHARE this checkout; **branch** from ``origin/master`` explicitly, re-verify HEAD before commit, stage explicit paths, serialize repo writes.**\*\*** The #116/#117 collision (owner upset). Global form in `~/.claude/CLAUDE.md`.

13 - [Gotcha: long GPU/WSL runs](gotcha-long-runs-gpu-wsl.md) ? running the WASB pipeline E2E on a big 4K video: Modern Standby kills it (keep-awake), cv2 extraction is the bottleneck (ffmpeg pre-extract), proxy-index/4K-stitch recipe + how to observe each phase.

14 - [Cache-hygiene policy](cache-hygiene-policy.md) ? retention table (purge frames/staged vs keep CSVs/DB/reels) + the shipped tooling: self-cleaning runners (keep\_frames default false), tools/cleanup\_caches.py (dry-run-first), disk guard, SessionStart nudge hook. The ``for e in *`` WSL footgun.

15 - [Tool: trim\_video.py](tool-trim-video.md) ? reusable ``scratch/trim_video.py`` for keyframe-aware LOSSLESS video trims (drop first/last N sec), ``--exact`` re-encode with NVENC?libx264 fallback. Reach for it instead of hand-rolling ffmpeg trims.

#### 567. user (2026-06-14T06:32:21.148Z)

The file `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md` has been updated successfully. (file state is current in your context ? no need to Read it back)

#### 568. assistant (2026-06-14T06:32:41.172Z)

Both done ? the convention is now written in **docs** and **memory**, and the run-drivers are segregated.

### What landed ([#138](https://github.com/avidullu/badminton-highlight-indexer/pull/138), merged)

- **Moved** ``run_phase3_eval.py`` + ``run_process_and_stitch.py`` ? **scripts/** (the two with zero importers). Updated the ``ruff.toml`` E402 paths + the doc/docstring references (1 line each, no behavior change).
- **Grandfathered** at root (referenced by CI/packaging/tests): ``setup.py``, ``run_all_tests.py``, ``run_regression.py``, ``run_eval.py``.
- **Left** ``backend/eval/`` alone ? those modules *look* like scripts but are **library** (``calibrate_wasb`` alone is imported by 8 sites); moving them would break the import graph for no gain.

### The guidelines you asked for ? in docs **and** memory

- **Docs** (authoritative): **new** ``docs/CODE_MAP.md`` §0 ? "REPO LAYOUT ? where new files go": a placement table (backend library vs ``scripts/`` vs ``tools/`` vs ``scratch/`` vs ``training/`` vs ``artifacts/`` vs ``output/`` ?) + the canonical-root-files list. Plus ``scripts/README.md``. `CODE_MAP` was the right home ? it's the cold-start map, and it had no placement section before (the gap you spotted).
- **Memory**: ``file-placement-conventions.md`` (type ``reference``) + a **MEMORY.md** pointer flagged "read before committing a new file" ? so future sessions (me or others) see it during ramp-up and route files correctly.

### One heads-up

The shared checkout is **currently** flipped to a sibling session's branch (``docs/i3-fusion-measured``), so your local view still shows the *old* root layout. That's just the local checkout ? ``origin/master`` is correctly tidied ( #138 merged at ``8836dfe``, with ``scripts/`` populated). It'll reflect once the checkout is back on master (a ``git checkout master && git pull`` when the sibling's done, or it'll sync naturally).

That wraps your cleanup asks: scratch gitignored ( #136), root run-drivers segregated + placement guidelines documented in docs + memory ( #138). The directory's tidy and the "where does this go?" question is now a 5-second lookup. ?

#### 569. user (2026-06-14T06:32:43.365Z)

please go ahead and merge the PRs from your wrok. Appreciate the diligence!

### 570. assistant (2026-06-14T06:33:00.343Z)

Let me verify the actual open-PR state ? I self-merged each of my PRs on green CI as I went (per the standing agreement), so I want to confirm whether anything of mine is still open before claiming it's all done.

### 571. user (2026-06-14T06:33:12.660Z)

```
=== ALL open PRs (number · branch · author · title) ===
#137 docs/i3-fusion-measured @avidullu ? docs(quality): log fusion P3+P4 measured ?F1 (n=6
golden ? honest negative/uncertain)
```

```
=== my merged PRs this session (mine = the ones I authored/landed) ===
```

```
#119: MERGED (2026-06-13T20:38:51Z)
#124: MERGED (2026-06-13T20:53:29Z)
#128: MERGED (2026-06-14T03:55:31Z)
#129: MERGED (2026-06-14T04:03:10Z)
#130: MERGED (2026-06-14T04:23:56Z)
#131: MERGED (2026-06-14T04:29:26Z)
#132: MERGED (2026-06-14T05:36:39Z)
#133: MERGED (2026-06-14T05:49:51Z)
#135: MERGED (2026-06-14T05:58:04Z)
#136: MERGED (2026-06-14T06:20:38Z)
#138: MERGED (2026-06-14T06:31:04Z)
```

```
=== #116 (the one that was closed-not-merged) ===
CLOSED
```

### 572. assistant (2026-06-14T06:33:43.870Z)

Good news: **there's nothing left to merge from my work** ? all 11 of my PRs this session are already merged. **I self-merged each on green CI as I went (per your standing merge-on-verified-green authorization), so they landed in real time rather than queuing up.**

The ledger (all `MERGED`):

```
| Track | PRs |
|---|---|
| Analyzers/reconcilers framework + lit review | #119 |
| Eval honesty (C1?C4 + wiring + learned fix) | #124, #129, #130 |
| First A/B experiment + Gate-A promotion | #128, #131 |
| Execution-policy (design + P1 + P2) | #132, #133, #135 |
| Repo hygiene (gitignore scratch + scripts/ + layout guidelines) | #136, #138 |
```

(#116 is the only one not merged ? it was **closed** because its content had already landed via the sibling's #117.)

### One open PR ? but it's **\*not\*** mine

The only thing open is **[#137](https://github.com/avidullu/badminton-highlight-indexer/pull/137)** `docs(quality): log fusion P3+P4 measured ?F1 (honest negative/uncertain)` ? that's a **sibling session's** fusion-measurement work (the GPU ?F1 run the handoff mentioned), a different track from my execution-policy/analyzers/eval work. I've deliberately **not** merged it: it's not my work, I don't have its context, and merging another session's in-flight PR blind would be presumptuous (and would mean vouching for numbers I didn't produce).

If you **do** want it landed, just say so and I'll check its CI and merge it ? but I'd rather you point me at it explicitly than assume. Otherwise it's the sibling's (or your) call.

So you're all set: everything I built is on `master`, the tree's tidy, and the next session has a clean P3 pickup. Truly enjoyed this one ? thank you for the trust and the great ideas. ?

### 573. user (2026-06-14T06:37:46.022Z)

please d merge it and thanks for the work!



#### 574. assistant (2026-06-14T06:37:59.915Z)

You've got it ? you've explicitly asked, so I'll merge #137. Verifying its CI is green and it's mergeable first (the same bar I held my own to), then merging.

#### 575. user (2026-06-14T06:38:03.333Z)

```
=== #137 mergeable + checks ===
{"files":["docs/QUALITY_ITERATIONS.md"],"mergeStateStatus":"CLEAN","mergeable":"MERGEABLE"}
Full test suite (offline, mocked)      pass    1m31s  https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490668085/job/81255002118
Mypy (lenient baseline) pass      31s      https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490668085/job/81255002100
Ruff lint          pass     9s      https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490668085/job/81255002103
Syntax sweep + import smoke (no GPU stack) pass    18s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27490668085/job/81255002119
```

#### 576. assistant (2026-06-14T06:38:11.785Z)

All green, mergeable, docs-only (the fusion P3/P4 ?F1 log entry). Merging.

#### 577. user (2026-06-14T06:38:17.563Z)

```
=== confirm ===
{"mergedAt":"2026-06-14T06:38:16Z","state":"MERGED"}
=== open PRs now ===
0
```

#### 578. assistant (2026-06-14T06:38:35.582Z)

Merged ? ? \*\*[#137](https://github.com/avidullu/badminton-highlight-indexer/pull/137) is in\*\* (the fusion P3/P4 measured-?F1 quality log, docs-only, all CI green), and there are now \*\*0 open PRs\*\* on the repo.

So everything's landed: my 11 PRs from this session plus the sibling's #137. `master` is clean, the tree's tidy, the handoff + resume prompt are ready, and the file-placement guidelines are in docs + memory.

(Tiny note for when you're back: your local checkout is still sitting on the now-merged `docs/i3-fusion-measured` branch ? a `git checkout master && git pull` will bring it current. I left it alone since it's the shared checkout.)

Genuinely a great session ? the prior-art review, the honest A/B, Gate-A, the policy framework. Thanks for the trust, the sharp ideas, and for caring about doing it \*right\*. Catch you in the next one. ?

#### 579. user (2026-06-14T06:38:38.056Z)

do merge it

#### 580. user (2026-06-14T06:38:55.455Z)

```
=== open PRs right now ===
(empty above = none open)
=== #137 state (the one just merged) ===
{"mergedAt":"2026-06-14T06:38:16Z","state":"MERGED"}
```

#### 581. assistant (2026-06-14T06:39:04.122Z)

It's already done ? \*\*#137 was merged a moment ago\*\* (at `06:38 UTC`), and a fresh check shows \*\*0 open PRs\*\* on the repo. Nothing left to merge (your UI was probably just showing a stale "merge" button before it refreshed).

Final state: everything's on `master` ? my 11 PRs from this session + the sibling's #137 ? and the open-PR count is zero. You're fully wrapped up. ? Thanks again, and see you next session! ?