

# claude-windows / 8d4866bf-8edd-4da9-909e-022ddee12675

## Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-inde  
xer\8d4866bf-8edd-4da9-909e-022ddee12675\subagents\workflows\wf_f9255361-69d\agent-  
a19ab9c03f2c58d9a.jsonl`  
- SHA-256: `ae706c7f9cb34e29910b9e7ba305e9932de698d9548f3658f854b55a37589688`  
- Source modified: `2026-06-22T09:35:56+00:00`  
- Imported at: `2026-07-05T16:48:25+00:00`  
- project: `wf_f9255361-69d`  
- session_id: `8d4866bf-8edd-4da9-909e-022ddee12675`
```

## Transcript

### 1. user (2026-06-22T09:32:46.023Z)

You are auditing ONE doc in the KhelSutra brand/site repo for staleness, as part of a "are all the docs up to date?" sweep.

Repo root: C:\Users\avidu\Projects\kheIsutra (bash path: /c/Users/avidu/Projects/kheIsutra)  
The repo is on `master`, freshly synced to remote HEAD ? the latest merged PR is #94. Today is 2026-06-22.

DOC TO AUDIT: docs/ANNOTATE\_AND\_TRAIN.md

Your job: read the doc IN FULL, extract every concrete, falsifiable claim, and VERIFY each against the LIVE repo ? do NOT trust the doc's own wording. Hard discipline rule: a claim is "correct" ONLY when you have file:line or command-output evidence for it; if you cannot verify it, mark verdict "unverifiable" ? never guess, never assume the doc is right. Re-check against actual code/git.

Kinds of claims to check:

- File / path / route references ("served at /capture", "see src/store.js", "public/og.png", "systemd kheIsutra-site") -> does the path or route ACTUALLY exist now? Confirm with git ls-files / Read / the routing in site/server.js + site/src/index.js.
- PR numbers and counts ("PRs #1-#8 merged", "20 tests green", "6 languages", "3-step wizard", "P0..P5 done") -> verify against `git -C <root> log`, the actual test output/count, and the code.
- Status markers (DONE / pending / IN PROGRESS / DRAFT / "?" / "owner-gated" / "NEXT") -> still accurate vs the code & git history? Is something marked pending that is actually shipped, or vice-versa?
- Dates, "live at <url>", "merged in #N".
- Whether the doc describes something now SUPERSEDED or fully shipped -> is it an archive candidate (a tracker whose work is complete, or content duplicated by docs/archives/)? Note: docs/archives/FRIENDLY\_AND\_MULTILINGUAL.md already exists as the archived multilingual tracker.

Tools: Read the doc + every cited file; Grep/Glob to locate code; Bash for `git -C /c/Users/avidu/Projects/kheIsutra log --oneline`, `... ls-files`, and greps. To check a test count, you may run the suite (`cd /c/Users/avidu/Projects/kheIsutra/site && node --test` or `npm test`), but prefer cheap checks. Verify against master as it stands NOW.

Severity: high = a factually wrong claim a reader would act on (wrong path/route, a "DONE" that isn't, a wrong count/PR that misleads). medium = outdated-but-harmless (stale PR range, old "where we are", superseded framing). low = cosmetic.

Do NOT modify any file ? audit only. Return ONLY the structured object (doc, overallStatus, findings[], summary).

### 2. assistant (2026-06-22T09:32:48.199Z)

I'll start by reading the doc in full, then verify its claims against the live repo.

### 3. user (2026-06-22T09:32:49.623Z)

```
1      # Annotate your footage ? train your own model (KhelSutra Studio USP pages)
2
3      > **Status:** `DRAFT` ? owner-gated on scope . **Owner:** Avi . **Created:** 2026-06-20
  . **Last updated:** 2026-06-20
4      > **Lifecycle:** `DRAFT` ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
  DONE)
5      > **Tracking anchors:** $7 progress tracker is the source of truth; index in
  `docs/README.md`; pointer in `NEXT_STEPS.md` + memory `current-state` once work starts.
6      > **Relation:** peer of `khelsutra/docs/PER_RALLY_SELECTION.md` and
  `khelsutra/docs/LOCAL_APPLIANCE_ARCHITECTURE.md` (the latter lands with PR #13); mirrors the
  engine repo's `PROJECT_DOC_TEMPLATE` discipline. Documents an EXTERNAL public tool (`rally-
  annotator`, MIT) and the engine's `backend/worker train` CLI; it does not own those repos.
7      > **Honesty note:** every engine-fact claim is tagged `[verified]` (read against code
  this session), `[design]` (proposed here), or `[researched]` (needs sign-off). Not legal advice.
8
9      ---
10
11     ## 0. TL;DR
12     Ship ONE new static **marketing page** (`site/public/own-model.html`, slug `/own-model`)
  with three anchored sections ? **Annotate your footage . Train your own model . Setup & use** ?
  telling the bring-your-own-model story and linking the real, public, MIT **rally-annotator** VLC
  plugin `[verified]`. The single most important caveat: only the **ANNOTATE half is real today**
  ? the rally-annotator ships (v1.6.4, MIT) and its CSV is directly ingestible by the engine with
  **no transform** `[verified]` ? while the **TRAIN half is an open, in-build CLI** whose GPU-free
  path trains a **pipeline-validation mock only**, and whose **real model needs a GPU + an in-
  flight native WASB runner** behind an **owner-gated ship decision** `[verified gate_verdict
  harness.py:98]`. The Studio SPA gets a **design-tracked, owner-gated** correction/contribute
  screen, **never a "Train now" button**. Defensible claim = **ownership / licensing / privacy /
  consent / minors-excluded**, NOT accuracy-superiority (there is no benchmark vs a paid LLM).
13
14     ## 1. Problem & goal
15     KhelSutra's differentiator vs paid LLMs is that **you can own a model trained on YOUR
  footage** ? your courts, your camera angles, your busy multi-court halls ? on a commercially-
  clean, inspectable stack. The concrete, usable entry point already exists: **rally-annotator**,
  a free MIT VLC Lua extension that turns "watch a match" into "produce golden labels" `[verified
  README.md:22-24]`. What's missing is the **education/positioning surface** that tells this story
  and the **honest documentation** of the end-to-end loop. "Good" = a viewer understands the moat,
  can install the annotator and produce labels today, and is **never misled** into thinking one-
  click training exists. Read first: `rally-annotator/README.md`, `rally-
  annotator/docs/CSV_FORMAT.md`, `badminton-highlight-indexer/backend/worker/__main__.py`,
  `khelsutra/docs/PER_RALLY_SELECTION.md`, `khelsutra/site/public/index.html` (the honesty voice).
16
17     ## 2. Decisions locked
18     | # | Decision | Source / date | Implication |
19     |---|-----|-----|-----|
20     | D1 | Marketing page ships NOW; SPA screen is design-tracked + owner-gated | This
  synthesis 2026-06-20 | Convert surface first; operate surface deferred |
21     | D2 | ONE marketing page, three anchored sections (Annotate/Train/Setup) | Avoid thin
  pages; annotator is external | Clone `capture.html` skeleton; no build step |
22     | D3 | Annotator docs live on marketing; SPA only links out | `[verified]` annotator is
  a standalone VLC tool | Do not re-author VLC instructions in the SPA |
23     | D4 | NO "Train now" button/endpoint in the SPA | `[verified]` training is offline
  `backend/worker train` | SPA's honest role = correction loop ? grow corpus |
24     | D5 | Frame USP as ownership/licensing/privacy/consent, NOT accuracy | `[verified]` no
  user-model-vs-LLM benchmark exists | No accuracy-superiority copy anywhere |
25     | D6 | Paid LLMs never a training source; MIT/Apache/BSD only; minors excluded; opt-in |
  `[verified CLAUDE.md]` + `harness.py:178,189-192` attestation | Reuse the site's existing minors
  + opt-in wording (`index.html:342-349`) |
26     | D7 | Annotator CSV needs NO transform for the engine label path | `[verified]` $4.2
  schema bridge | Document the MD5+stem join as the only real prereq |
27
```

```

28     ## 3. Foundation ? research / prior art `[verified this session]`
29     **The annotator (real, shipping).** `rally-annotator` is a single-file VLC **Lua
extension** (`vlc/rally_annotator.lua`, VERSION 1.6.4), MIT-licensed, public at
https://github.com/Khelsutra/rally-annotator. It is button-driven (no listener capabilities;
deliberate ? VLC pause-listener is flaky on macOS / broken in 4.0), reads playback time via
`vlc.var.get(input, 'time')` (micros/1e6?s, lines 201-207), and appends one row per rally to
`<video>.rallies.csv` next to the video with atomic rewrite (save_all 310-335). Targets 5
net-separated racquet sports sharing one error taxonomy. CI runs a Lua 5.1 dialog suite +
layout-snapshot golden diff. VLC **3.0.x only** (4.0 changed the Lua API).
30
31     **The train loop (partly real).** The engine's `backend/worker train` (cmd_train,
`__main__.py:231`) pulls `labels/` (+ `videos/` for the detector source), builds per-video
trajectory CSVs, assembles a manifest, **shells out** to `python -m training.gen0.harness` (the
one-way train/serve wall, `training/README.md:9-11`), and pushes a versioned bundle to
`weights/<version>/`. Two trajectory sources, identical downstream: **synth** (DEFAULT, GPU-
free, a CPU mock that ENCODES the labels ? plumbing-meaningful only,
`stub_runner.py:124-149,131-133`) and **detector** (real WASB, M3.5, **GPU-gated, IN-FLIGHT** ?
the Linux-native runner does not exist yet, `02:15`, `04:98-113`). The harness fits a ~1M-param
LogisticRegression on CPU and emits the artifact bundle (`weights.json` + reviewable
`manifest.json`); **serving consumes only the bundle, never a code import** (ADR-008). The
**ship-gate** (`gate_verdict`, `harness.py:98`) currently refuses a commercial ship** regardless
of F1: WASB checkpoint provenance C3 unresolved (commercial ship refused) + F1@tIoU not yet
calibrated (#172).
32
33     **Guardrails (verbatim, non-negotiable).** Paid LLMs (Gemini/OpenAI/Anthropic) =
inference/serving/fallback ONLY, **never a training data source** (bundle attests
`no_paid_llm_training_labels:True`, `harness.py:192`); label source must be human (annotator
golden CSVs, `harness.py:178`). Commercial-clean deps only ? **MIT/Apache-2.0/BSD**. Owned base
= Qwen3-VL (Apache) primary; Gemma-4 AMBER/bench-only; Gemma 3 rejected. **Minors excluded**
from the training pool; per-user training data needs **explicit opt-in** `[verified CLAUDE.md]`.
34
35     ## 4. Design / architecture
36
37     ### 4.1 The annotate ? CSV ? train ? model flow `[verified except where noted]`
38     1. **ANNOTATE** (rally-annotator, VLC): per rally, Mark START (serve-contact) ? Mark END
(shuttle-dead) ? Ending reason ? optional shots ? Save Rally ? appends a row to
`<video>.rallies.csv`. No compute, no GPU, no LLM.
39     2. **STAGE**: `python -m src.cli.curate import-local --sport badminton --video-path
match.mp4 --annotations-path match.rallies.csv`, and/or populate a corpus prefix `videos/
labels/ manifest.json weights/ outputs/<video_id>/`.
40     3. **TRAIN**: default `python -m backend.worker train --corpus-uri ? --out-uri ?
--version 0.1.0` (synth, GPU-free CPU mock, ?2 labeled videos); real = add `--trajectory-source
detector --segmenter wasb_hybrid` with `indexing.detector_impl=native` (GPU + native WASB +
`wasb_badminton_best.pth.tar`, IN-FLIGHT).
41     4. **HARNESS ? BUNDLE**: parse trajectory + load GT ? frame features ? honest LOVO eval
? fit final LR + tune threshold ? ship-gate ? emit `weights.json` + `manifest.json`.
42     5. **USE** `[researched/design, owner-gated]`: if Gen-0 clears the ship-gate A/B, the
bundle drives the existing record?upload?reel loop. No per-user serving UI; no Studio Train
button.
43
44     ### 4.2 The schema bridge ? annotator CSV ? engine label CSV (MATCHES, no transform)
45     Annotator writes header
`rally_number,start_time,end_time,ending_reason,sport,shots_count` (`rally_annotator.lua:128`;
`CSV_FORMAT.md:6`), `%.3f` decimal seconds, end>start. The engine reads **by column name** via
two readers, both satisfied:
46     - **Default synth path** ? `synth_trajectory_from_labels` uses `csv.DictReader` reading
literally `row["start_time"]`/`row["end_time"]` (`stub_runner.py:139`) ? **MATCHES exactly** ;
extra columns ignored. (This path needs the `_time` names specifically ? the short `start/end`
alias is NOT accepted here.)
47     - **Canonical GT reader** ? `load_gt_intervals` accepts `start`/`end` OR
`start_time`/`end_time`, ignores other columns, also accepts headerless positional col0,col1,
tolerates decimal-seconds or MM:SS/HH:MM:SS, BOM, blank/`#` lines (`gt_loader.py:14-90`) ?
**MATCHES** the `_time` branch.
48
49     The worker's own inference CSV uses `start,end,shot_count,ending_reason`

```

(`\_\_main\_\_.py:62-74`) ? order/short-name differ, but order never matters and `gt\_loader` accepts both aliases; the annotator's `\*\_time` form is the **\*\*safer\*\*** of the two. **\*\*The ONE** real prerequisite is an IDENTITY/JOIN constraint, not a schema gap:**\*\*** training keys video?label by **\*\*file MD5\*\*** and stem-matches `labels/<name>.rallies.csv` ? `videos/<name>.<ext>` (``__main__.py:271-280``; ``02:48-60``). Stems must agree and the staged video bytes must be **\*\*byte-identical\*\*** to what was annotated ? any re-encode changes the MD5 and breaks the join (collector manifest currently omits `md5`, M2 prereq, ``02:62-68``).

50  
51     **### 4.3 Reuse map**  
52     - **\*\*Marketing\*\*** `[verified]`: clone `site/public/capture.html` skeleton + brand tokens (``index.html:25-32``); reuse ``.honest`/`.notes`` components (``:152-158``, used ``:351-354``, ``:400-415``); add a ``.own-model`` alias in ``site/server.js:49-51`` (Worker ``.env.ASSETS.fetch`` needs no change, ``site/src/index.js:84``); cross-link nav ``:218-223``, the "gets sharper" section ``:327-356``, footer ``:528``; screenshots under ``site/public/img/`` (precedent ``site/og.png`` + ``.make_og.py``).

53     - **\*\*SPA\*\*** `[design, gated]`: today ``web/src/App.tsx`` is one static screen, **\*\*no router\*\*** (``web/package.json``). A correction screen needs react-router-dom (or a hash switch) + the typed client ``web/src/api/client.ts`/`types.ts`` + ``web/src/lib/log.ts``; it realizes P10 in ``PER_RALLY_SELECTION.md:122`` and depends on a **\*\*new\*\*** engine endpoint ``PATCH /api/segments/{id}`` (dead columns ``candidate_segments.human_label/notes/confirmed_segment_id``).

54     - **\*\*Annotator content\*\*** `[verified]`: embed CSV example + install commands + ending-reason table **\*\*verbatim\*\*** from the repo; link, don't re-author.

55  
56     **## 5. Threat model / privacy**

57     - **\*\*Footage stays the user's\*\***: the annotator runs locally in the user's own VLC and writes a CSV next to the video ? labeling happens with **\*\*footage never leaving the machine\*\*** `[verified]`.

58     - **\*\*Training use = explicit opt-in, never automatic\*\*** `[verified CLAUDE.md]`; reuse the site's existing opt-in wording.

59     - **\*\*Minors excluded from any training pool\*\*** `[verified CLAUDE.md + live site "Minors are protected"]` ? keep prominent (audience = academies with children).

60     - **\*\*No paid-LLM laundering\*\***: the model never trains on a paid LLM's output (terms/copyright), attested in the bundle (``harness.py:192``).

61     - **\*\*Residual\*\***: base-model (Qwen3-VL) pretraining provenance is unauditabile ``.researched`` ? owner risk-acceptance required before VLM "train your own model" copy goes beyond roadmap.

62  
63     **## 6. Honest limits ? what this does NOT do**

64     A shipped marketing page does NOT mean training is one-click (it is an offline CLI), does NOT mean a real model exists (synth = mock; real = GPU + in-flight native WASB + owner-gated ship-gate), and does NOT imply any accuracy claim vs a paid LLM (no benchmark exists). The annotator is a **\*\*separate VLC 3.0.x desktop plugin\*\***, not in-browser; it is **\*\*click-driven\*\*** (no marking hotkeys). The SPA correction screen is **\*\*design-tracked + owner-gated\*\*** and depends on an unbuilt ``PATCH /api/segments`` endpoint. Two CSV schemas exist (annotator lighter vs Roora vendor) ? do not conflate.

65  
66     **## 7. Deliverables & progress tracker ? source of truth**

67     Legend: ? Todo · ? In progress · ? Done · ? Gated. One small PR per row, branched from ``origin/master``, no stacking.

68  
69     | ID | Deliverable | Depends on | Gated? | Status | PR |  
70     |----|-----|-----|-----|-----|----|  
71     | P0 | This tracked doc in ``khelsutra/docs/``, indexed in ``docs/README.md`` | ? | No | ? |  
72     | P1 | ``site/public/own-model.html`` (clone ``capture.html``; 3 sections; honesty fences; placeholder shots) | P0 | No | ? | (this PR) |  
73     | P2 | ``.own-model`` alias in ``site/server.js`` (Worker needs no change) | P1 | No | ? | (this PR) |  
74     | P3 | Cross-links: nav + footer (the "gets sharper" section link is a nice-to-have, deferred) | P1 | No | ? | (this PR) |  
75     | P4 | Embed verbatim repo text (CSV/install/ending-reasons) + GitHub link | P1 | No | ? | (this PR) |  
76     | P5 | Capture + commit real screenshots (consented, no minors, faces redacted); replace placeholders | P1 | ? Owner (footage) | ? | ? |  
77     | P6 | Owner: slug + 1-vs-3-page + train-positioning honesty | ? | ? Owner | ? | ? |

```

78      | P7 | [design] SPA correction/contribute route (RallyReview + link-out) | P8 | ? Owner
| ? | ? |
79      | P8 | [design] Engine `PATCH /api/segments/{id}` (M4 corpus extension) | ? | ? Owner
(engine) | ? | ? |
80      | P9 | [researched] Serve a shipped Gen-0 bundle into record?reel | native WASB + ship-
gate | ? Owner+blocked | ? | ? |
81
82      ## 8. Open questions ? owner / external
83      **Blocking gates:** (1) Greenlight how prominently to position "train your own model"
(CLI + roadmap). (2) Approve the SPA correction screen + the new `PATCH /api/segments/{id}`
endpoint. (3) Owner risk-acceptance on Qwen3-VL provenance before VLM copy. (4) Provide
consented no-minors footage for the 5 screenshots.
84      **Nice-to-knows:** slug (`/own-model` vs `/byo-model` vs `/annotate`); one page vs
three; react-router-dom vs hash switch; cross-link `docs/LOCAL_APPLIANCE_ARCHITECTURE.md` once
PR #13 merges.
85
86      ## 9. Definition of done
87      - [ ] `own-model.html` live with Annotate/Train/Setup sections, honesty fences, verbatim
repo text, GitHub link, real (non-fabricated) screenshots.
88      - [ ] `own-model` alias works on both the Node host and the Cloudflare Worker; cross-
linked from nav/section/footer.
89      - [ ] Every claim tagged `[verified]`/[design]`/[researched]`; no accuracy-superiority
claim; no "Train now" button; VLC-3.0.x-only + GPU/native-WASB-gated + owner-gated stated
plainly.
90      - [ ] Schema-bridge + MD5/stem-join prerequisite documented; the two CSV schemas kept
distinct.
91      - [ ] $7 tracker reflects real shipped state; doc indexed in `docs/README.md`.
92
93      ## 10. References
94      **Annotator `[verified]`:** `rally-annotator/README.md`, `vlc/rally_annotator.lua`
(HEADER :128; now_seconds :201-207; marks :608-628; save_all :310-335; resolve_out_path
:253-267), `docs/CSV_FORMAT.md`, `docs/ENDING_REASONS.md`, `LICENSE` (MIT),
`.github/workflows/ci.yml`.
95      **Engine `[verified]`:** `backend/worker/__main__.py` (cmd_train :231;
_write_intervals_csv :62-74; join :271-280; harness shell-out :321-328),
`backend/pipeline/detectors/stub_runner.py` (synth :124-149, :131-133, :139),
`training/gen0/harness.py` (gate_verdict :98; label_source :178; attestation :189-192),
`backend/eval/gt_loader.py` (:14-90), `docs/SETUP_NEW_MACHINE.md`,
`docs/DECOUPLED_COMPUTE/02-COMPUTE-MAP-AND-CONTRACT.md`, `docs/DECOUPLED_COMPUTE/04-BUILD-AND-
TEST-PLAN.md`, `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`, `CLAUDE.md`.
96      **KheISutra `[verified]`:** `site/public/index.html`, `site/public/capture.html`,
`site/server.js`, `site/src/index.js`, `web/src/App.tsx`, `web/package.json`, `docs/README.md`,
`docs/PER_RALLY_SELECTION.md`.
97
98      ### Changelog
99      - 2026-06-20 ? Initial DRAFT consolidating annotator + train-loop + UI-surface research
into one tracked doc (multi-agent design pass; schema-bridge + no-images + train-guardrails re-
verified against source before writing).
100

```

#### 4. user (2026-06-22T09:32:50.811Z)

```

50b2bd7 chore(repo): ignore local runtime artifacts + correct stale NEXT_STEPS claims (#94)
e2570e4 docs: archive the Friendly & Multilingual tracker (engineering complete) (#93)
0bef61e docs(friendly-multilingual): close out the chapter ? log #89/#90/#91 + strengthen fitNav
tests (#92)
9f77ddc test(site): browser layout-regression CI guard + auto-collapse the nav to fit every
locale (#91)
4415fe8 feat(site): localize remaining home sections (6 langs) + smooth the responsive header
(#90)
26c9da1 feat(site): localize the hero demo mockup + #how feature block (6 languages) (#89)
7ae9cee feat(web): P2c first-run setup wizard (3 steps from /api/capabilities, 6 languages)
(#88)
cf59fbf feat(site): add es/da/id to the marketing site (parity with the SPA's 6 languages) (#87)
17242a9 feat(site): localize own-model body content to kn/hi (#86)

```

```

7b04f02 deploy(kit): align with the live-verified droplet deploy (+ the gemini-default fix)
(#85)
5e5a2ac docs(friendly-multilingual): close out the chapter (§9 DoD + UI/UX + B7 deferred) (#84)
4c0f35c fix(site): UI/UX ? compact language switcher + header overflow fix + localized hero
(#83)
0fae364 feat(web+site): A2 ? localized WhatsApp share for the reel (SPA + site) (#82)
1e52cfd deploy(droplet): tailored alpha kit for the CPU droplet (engine api. side) (#81)
64076a5 test(web): add a coverage floor (CI-gated) + i18n regression tests for the #77 audit
(#80)
---BRANCH---
master
---HEAD---
50b2bd7 chore(repo): ignore local runtime artifacts + correct stale NEXT_STEPS claims (#94)

```

## 5. assistant (2026-06-22T09:32:55.066Z)

Now let me verify the claims. The doc claims P1-P4 are "(this PR)" done and references many files. Let me check the actual files exist and verify the engine/annotator claims. Let me start with the KhelSutra repo files since those are directly checkable.

## 6. user (2026-06-22T09:32:57.480Z)

```

=== own-model.html exists? ===
site/public/own-model.html
=== capture.html exists? ===
site/public/capture.html
=== server.js own-model alias ===
41:const LOCALIZED_PAGES = new Set(["index.html", "capture.html", "own-model.html"]);
99:    : pathname === "/own-model" ? "own-model.html"
=== src/index.js ASSETS ===
84:    return env.ASSETS.fetch(request); // everything else: static assets

```

## 7. user (2026-06-22T09:32:59.493Z)

```

=== PRs touching own-model.html ===
cf59fbf feat(site): add es/da/id to the marketing site (parity with the SPA's 6 languages) (#87)
17242a9 feat(site): localize own-model body content to kn/hi (#86)
4c0f35c fix(site): UI/UX ? compact language switcher + header overflow fix + localized hero
(#83)
1cd84af feat(site): A3 ? localize recording-education + marketing to kn/hi, Noto on the site
(#74)
1b1fa7b feat(web+site): B3 ? trilingual marketing site via shared t()/data-i18n shim (D-CATALOG-
HOME) (#70)
26d636d feat(site): link the academy flyer from the top nav + own-model page (#44)
9f0942b feat(site): own-model ? add an obvious back-to-home link; backlog: sister-repo org move
(#26)
1906c82 feat(site): own-model ? gate the reel, show the real annotator screenshots + a "default
works" note (#25)
f09e348 feat(site): reuse the sample reel on the own-model page (click-to-play, privacy-
consistent) (#24)
87520b8 chore(site): point rally-annotator links at the new Khelsutra org (#23)
b1110e1 feat(site): own-model page ? annotate your footage & train your own model (annotate-
train P1-P4) (#22)
=== PR that added doc P0 ===
87520b8 chore(site): point rally-annotator links at the new Khelsutra org (#23)
b1110e1 feat(site): own-model page ? annotate your footage & train your own model (annotate-
train P1-P4) (#22)
6b62ebf docs(annotate-train): plan rally-annotator integration + "train your own model" USP
pages (#15)
=== docs/README.md mentions ANNOTATE_AND_TRAIN? ===

```

## 8. assistant (2026-06-22T09:33:03.360Z)

The doc was created in #15 (P0), then own-model.html shipped in #22 (P1-P4), with screenshots in

#25. Let me check the doc index claim (docs/README.md has no mention), the server.js alias details, and whether P5 (screenshots) actually shipped despite being marked ?. Also let me verify the localized-pages and screenshot facts.

## 9. user (2026-06-22T09:33:05.676Z)

=== docs/README.md full content ===

# khelsutra docs index

## Tracked work (this repo)

- `**`DRY_RUN_CUJ.md`**` ? ? `**DELIVERED (2026-06-21):**` the UI drives the tool on a `**GCS-bucket video**` ? review/select ? compile ? `**reel downloaded local**`, verified in en + Hindi. Path A (local) + Path B/B-P2 (in-UI ``gs://`` ingest, [#56]) both live e2e with a screencast (\$9); the only remaining box (droplet, option 3) is `**parked by the owner**` (GPU subscription). Two engine bugs the runs surfaced are fixed (#281, #289). Was a milestone of ``LOCAL_APPLIANCE_ARCHITECTURE.md``.
- `**`archives/FRIENDLY_AND_MULTILINGUAL.md`**` ? ? `**DONE ? ARCHIVED (2026-06-22).**` Made the product super user-friendly for non-tech-savvy coaches/parents `**and**` localized it to all `**6 languages**` across both surfaces (``web/`` SPA ? react-i18next, ``site/`` ? a shared ``t()``` shim); the home-page body is fully localized (except the deferred ``#join`` form, A6) and the responsive header is fixed + CI-guarded. Implementation project for the engine's ``I18N_PLAN.md`` (l10n) + the UX half of ``UI_PROPOSAL.md``. `**Carried forward to `NEXT_STEPS.md`:**` native-speaker review + owner font-glyph eyeball (human gates), and `**B7**` (engine backend l10n, deferred).
- `**`MULTILINGUAL_EXPANSION.md`**` ? proposal to add `**7 more locales**` (th/id/da/es/zh-Hans/ja/ko) with en/kn/hi-level diligence; peer of ``FRIENDLY_AND_MULTILINGUAL.md``, reusing its foundation. ``DRAFT`` ? `**owner-gated (approval to start)**`, [#58].
- `**`LOCAL_APPLIANCE_ARCHITECTURE.md`**` ? the single-box, UI-powered rally indexer + reel generator (KhelSutra Studio): the ``web/`` SPA drives the full ingest ? index ? reel loop on one box behind an edge gate; the single-box deployment profile of the engine's decoupled architecture. Umbrella for the console; per-rally selection is one screen of it. ``DRAFT``.
- `**`PER_RALLY_SELECTION.md`**` ? interactive pick-rallies ? reel builder (RallyPicker). The alpha SPA over the engine's rally index; hybrid grid + deterministic query bar on the `**existing**` ``/api/query`` + ``/api/compile`` contract. ``IN PROGRESS`` (decisions locked).
- `**`REQUEST_ACCESS_FORM.md`**` ? segmented lead-capture form + Worker/D1 backend (shipped).
- `**`PORTABILITY.md`**` ? the Node+SQLite VM-host runbook (Cloudflare ? droplet parity).

## Engine planning docs

Canonical planning docs currently live in the engine repo (``badminton-highlight-indexer/docs``): ``UI_PROPOSAL.md`` · ``HOSTING_PLAN.md`` · ``UI_ALPHA_EXECUTION_PLAN.md``. They migrate here when the control plane lands (per ADR-009 in the engine repo's ``docs/archives/decisions/DECISIONS.md``).  
=== does docs/README.md exist? ===  
`docs/README.md`

## 10. user (2026-06-22T09:33:06.856Z)

=== screenshots in site/public/img ===

`site/public/img/annotate-enable.png`  
`site/public/img/annotate-tool.png`  
`site/public/img/flyer-qr.svg`  
`site/public/img/studio-en.png`  
`site/public/img/studio-hi.png`  
`site/public/img/studio-kn.png`

=== own-model.html screenshot refs ===

7:<meta name="description" content="With KhelSutra you can annotate your own footage and grow a rally model that's YOURS ? on an open, MIT-licensed stack you control. The free VLC rally-annotator turns 'watch a match' into golden labels." />

```

69:  /* annotation-tool screenshots (no footage/people ? pure UI) */
108:    <p data-i18n-html="site.ownModel.annotate.body"><b>rally-annotator</b> is a free,
open-source (MIT) <b>VLC plugin</b> that turns "watch a match" into golden labels. While you
watch your own video in VLC, you mark each rally's start and end plus how the point ended; it
writes one row per rally to a plain CSV next to your video. <b>Your footage never leaves your
machine.</b> It covers five net-separated racquet sports (badminton, tennis, table tennis,
pickleball, padel).</p>
109:    <a class="ghbtn" href="https://github.com/Khelsutra/rally-annotator"><svg width="16"
height="16" viewBox="0 0 16 16" fill="currentColor" aria-hidden="true"><path d="M8 0C3.58 0 0
3.58 0 8a8 8 0 0 0 5.47 7.59c.4.07.55-.17.55-.38 0-.19-.01-.82-.01-1.49-2.01.37-2.53-.49-2.69-
.94-.09-.23-.48-.94-.82-1.13-.28-.15-.68-.52-.01-.53.63-.01 1.08.58 1.23.82.72 1.21 1.87.87
2.33.66.07-.52.28-.87.51-1.07-1.78-.2-3.64-.89-3.64-3.95
0-.87.31-1.59.82-2.15-.08-.2-.36-1.02.08-2.12 0 0 .67-.21 2.2.82a7.4 7.4 0 0 1 2-.27c.68 0
1.36.09 2 .27 1.53-1.04 2.2-.82 2.2-.82.44 1.1.16 1.92.08 2.12.51.56.82 1.27.82 2.15 0 3.07-1.87
3.75-3.65 3.95.29.25.54.73.54 1.48 0 1.07-.01 1.93-.01 2.2 0 .21.15.46.55.38A8 8 0 0 0 16
8c0-4.42-3.58-8-8-8Z"/></svg> github.com/Khelsutra/rally-annotator</a>
111:    <figure><figcaption
data-i18n="site.ownModel.annotate.cap1">Enable it once: VLC ? View ? Rally
Annotator.</figcaption></figure>
112:    <figure><figcaption
data-i18n="site.ownModel.annotate.cap2">Mark each rally's start/end + ending reason, then Save
Rally.</figcaption></figure>
154:    <p data-i18n-html="site.ownModel.setup.install"><b>Install the plugin (VLC 3.0.x).</b>
Copy <code>rally_annotator.lua</code> into VLC's per-user Lua <code>extensions</code> folder
(create it if missing), then VLC ? Tools ? Plugins and extensions ? <i>Reload extensions</i>,
and enable it from <i>View ? Rally Annotator</i>.</p>
191:    <span><a href="/">khelsutra.guru</a> · <a href="https://github.com/Khelsutra/rally-
annotator">rally-annotator (MIT)</a> · <span data-i18n="site.ownModel.footer.flyer">Academy
flyer:</span> <a href="/flyer" hreflang="en">English</a> · <a href="/flyer-kn" hreflang="kn"
lang="kn">?????</a> · <a href="/flyer-hi" hreflang="hi" lang="hi">?????</a></span>

```

## 11. assistant (2026-06-22T09:33:15.741Z)

Several important findings emerging. Key ones:

1. **\*\*docs/README.md does NOT mention ANNOTATE\_AND\_TRAIN.md\*\*** ? the doc's tracking claim (line 5: "index in `docs/README.md`") and DoD ("doc indexed in `docs/README.md`") and \$7 P0 PR #15 claim "indexed in `docs/README.md`" are unverified ? the index is missing it.
2. **\*\*Real screenshots SHIPPED\*\*** (`annotate-enable.png`, `annotate-tool.png`) in PR #25 ? but \$7 P5 ("Capture + commit real screenshots... replace placeholders") is marked ? (Todo). That's stale ? it appears done.
3. **\*\*rally-annotator links point at `github.com/Khelsutra/rally-annotator`\*\*** (a NEW org), not the doc's stated `https://github.com/Khelsutra/rally-annotator`... wait, the doc line 29 says `https://github.com/Khelsutra/rally-annotator`. Let me check ? actually the doc says `Khelsutra/rally-annotator`. Matches. Good.

Let me verify the server.js alias mechanism more precisely, the localization (the doc says nothing about localization but it's now a localized page), and check the engine repo claims. Let me look at server.js routing and check the engine files.

## 12. user (2026-06-22T09:33:17.229Z)

```

38    // shim negotiates ? see src/index.js.)
39    const SITE_LOCALES = ["en", "kn", "hi", "es", "da", "id"];
40    const LANG_COOKIE = "khelsutra.lang";
41    const LOCALIZED_PAGES = new Set(["index.html", "capture.html", "own-model.html"]);
42
43    const baseLang = (code) => (code ?
String(code).toLowerCase().split(";")[0].trim().split("-")[0] : "");
44
45    function cookieValue(header, name) {

```



```

46     const m = (header || "").match(new RegExp("(?:^|;\\s*)" + name.replace(/\./g, "\\.") +
"=([^\s]*)"));
47     return m ? decodeURIComponent(m[1]) : null;
48 }
49
50 function parseAcceptLanguage(header) {
51     return (header || "")
52         .split(",")
53         .map((part) => {
54             const [tag, ...params] = part.trim().split(";");
55             const q = params.map((p) => p.trim()).find((p) => p.startsWith("q="));
56             return { tag, q: q ? Number(q.slice(2)) : 1 };
57         })
58         .filter((x) => x.tag)
59         .sort((a, b) => b.q - a.q)
60         .map((x) => x.tag);
61 }
62
63 // persisted (cookie) -> ?lang= -> Accept-Language -> en. Anything unsupported is
skipped.
64 function negotiateLang(req, searchParams) {
65     const candidates = [
66         cookieValue(req.headers.cookie, LANG_COOKIE),
67         searchParams.get("lang"),
68         ...parseAcceptLanguage(req.headers["accept-language"]),
69     ];
70     for (const c of candidates) {
71         const b = baseLang(c);
72         if (b && SITE_LOCALES.includes(b)) return b;
73     }
74     return "en";
75 }
76
77 const setHtmlLang = (html, lang) => html.replace(/<html lang="[^"]*" /i, `<html
lang="${lang}" `);
78
79 async function toRequest(req) {
80     const url = `http://${req.headers.host || "localhost"}${req.url}`;
81     let body;
82     if (req.method !== "GET" && req.method !== "HEAD") {
83         const chunks = [];
84         for await (const c of req) chunks.push(c);
85         body = Buffer.concat(chunks);
86     }
87     return new Request(url, { method: req.method, headers: req.headers, body });
88 }
89
90 async function send(res, response) {
91     res.statusCode = response.status;
92     for (const [k, v] of response.headers) res.setHeader(k, v);
93     res.end(Buffer.from(await response.arrayBuffer()));
94 }
95
96 async function serveStatic(pathname, req, searchParams, res) {
97     const rel = pathname === "/" ? "index.html"
98         : pathname === "/capture" ? "capture.html"
99         : pathname === "/own-model" ? "own-model.html"
100         : pathname === "/flyer" ? "flyer.html"
101         : pathname === "/flyer-kn" ? "flyer-kn.html"
102         : pathname === "/flyer-hi" ? "flyer-hi.html"
103         : pathname.replace(/^\/+/, "");
104     const file = normalize(join(PUBLIC, rel));
105     if (!file.startsWith(PUBLIC)) { res.statusCode = 403; return res.end("forbidden"); }
106     try {
107         let data = await readFile(file);

```

### 13. user (2026-06-22T09:33:20.173Z)

=== screenshot PR ===

1906c82 feat(site): own-model ? gate the reel, show the real annotator screenshots + a "default works" note (#25)

=== are these placeholders or real? file sizes ===

```
-rw-r--r-- 1 avidu 197609 133712 Jun 20 20:47 site/public/img/annotate-enable.png
-rw-r--r-- 1 avidu 197609 153133 Jun 20 20:47 site/public/img/annotate-tool.png
```

### 14. assistant (2026-06-22T09:33:24.667Z)

The `own-model` alias exists in server.js:99. The screenshots are real (133KB, 153KB PNGs from PR #25, "show the real annotator screenshots"). So P5 (real screenshots) is effectively DONE but marked ?.

Now let me verify the Worker (src/index.js) handles `own-model` ? the doc claims "Worker `env.ASSETS.fetch` needs no change". Let me check how the Worker maps `own-model`. Also let me verify engine repo claims.

### 15. user (2026-06-22T09:33:26.093Z)

```
40     org_name: str(p.org_name, 200),
41     courts: COURTS.includes(p.courts) ? p.courts : null,
42     has_gpu: GPU.includes(p.has_gpu) ? p.has_gpu : null,
43     storage,
44     message: str(p.message, 4000),
45     consent: p.consent ? 1 : 0,
46     country: str(meta.country, 8),
47     user_agent: str(meta.userAgent, 400),
48   };
49 }
50
51 function json(obj, status = 200) {
52   return new Response(JSON.stringify(obj), { status, headers: { "content-type":
"application/json" } });
53 }
54
55 export async function handleRequest(request, store, notify) {
56   let body;
57   try {
58     body = await request.json();
59   } catch {
60     return json({ ok: false, error: "invalid json" }, 400);
61   }
62   if (isBot(body)) return json({ ok: true }); // silently drop bots
63   const v = validate(body);
64   if (!v.ok) return json({ ok: false, error: v.error }, 400);
65
66   const r = toRow(body, {
67     country: request.headers.get("cf-ipcountry"),
68     userAgent: request.headers.get("user-agent"),
69   });
70
71   await store.insertAccessRequest(r);
72   try { await notify?.notify(r); } catch { /* email is best-effort; never fail the
request */ }
73
74   return json({ ok: true });
75 }
76
77 export default {
78   async fetch(request, env) {
79     const url = new URL(request.url);
80     if (url.pathname === "/api/request-access") {
81       if (request.method !== "POST") return json({ ok: false, error: "method not
```

```

allowed" }, 405);
82         return handleRequest(request, d1Store(env.DB), cfEmailNotify(env.NOTIFY,
env.NOTIFY_FROM, env.NOTIFY_TO));
83     }
84     return env.ASSETS.fetch(request); // everything else: static assets
85 },
86 };
87

```

#### 16. user (2026-06-22T09:33:27.201Z)

=== own-model alias in src/index.js (Worker)? ===

#### 17. assistant (2026-06-22T09:33:33.376Z)

Confirmed: the Worker (src/index.js:84) does `env.ASSETS.fetch(request)` for everything non-API, and has no `own-model` special case ? matching the doc's claim that "Worker `env.ASSETS.fetch` needs no change". Though note: the Worker relies on the static-asset platform to resolve `own-model` ? `own-model.html`; the doc's DoD claims "alias works on both the Node host and the Cloudflare Worker". The Node host has the explicit alias; the Worker depends on ASSETS routing config. That's a design nuance but the doc's "needs no change" claim is consistent with the code.

Now let me verify the engine repo claims. The doc references many `badminton-highlight-indexer` files. Let me check the key ones.

#### 18. user (2026-06-22T09:33:37.294Z)

Already up to date.

```

=== cmd_train in __main__.py ===
263: #: Video extensions accepted under <corpus>/videos/ (mirrors the labels/ '.csv' filter so a
271:   The detector trajectory is frame-indexed; the harness maps frame->time via fps, so a
WRONG
332: def cmd_train(args: argparse.Namespace) -> int:
334:     to training.gen0.harness (backend must NOT import training) -> push the artifact bundle.
336:     Two trajectory sources (the train pipeline + harness are identical either way):
359:     label_uris = sorted(u for u in storage.list_uris(f"{corpus}/labels/",
config=storage_cfg)
363:         f"under {corpus}/labels/")
383:     fname = storage.parse_uri(uri).name # e.g.
GX020094_proxy.rallies.csv
384:     name = fname.replace(".rallies.csv", "").replace(".csv", "")
424:     cmd = [sys.executable, "-m", "training.gen0.harness",
433:     print(f"[worker] gen0 harness failed (rc={rc})")
472:     pt = sub.add_parser("train", help="Gen-0 train-from-bucket (labels [+videos] ->
trajectories -> harness)")
473:     pt.add_argument("--corpus-uri", required=True, help="prefix containing labels/ (+
videos/ for ")
490:     pt.set_defaults(func=cmd_train)

```

#### 19. user (2026-06-22T09:33:38.116Z)

```

=== gate_verdict in harness.py ===
142: def gate_verdict(eval_res: dict, floor: Optional[dict]) -> dict:
196:     verdict = gate_verdict(eval_res, floor)
222:     "label_source": "human (rally-annotator golden CSVs)",
236:     "no_paid_llm_training_labels": True,
=== synth in stub_runner.py ===
124: def synth_trajectory_from_labels(labels_csv: str, out_csv: str, *, fps: float = 30.0,
137:     for row in _csv.DictReader(f):
139:         rallies.append((float(row["start_time"]), float(row["end_time"])))
=== gt_loader ===
4:- ``annotations.load_annotations``: positional ``start,end`` (the collector's
6:- ``calibrate_wasb.load_golden_gts``: DictReader requiring ``start_time/end_time``

```

```

15:- header with ``start``/``end`` OR ``start_time``/``end_time`` columns (any other
30:_START_NAMES = ("start", "start_time")
35:     ""If `row` is a header naming start/end columns, return their indices, else None.""
44: def load_gt_intervals(csv_path: str, *, strict: bool = True) -> List[Interval]:
48:     (bad timestamp, end <= start, negative start) so labeling mistakes fail loudly.
52:     Returns intervals sorted by start; overlapping GT intervals are logged as a warning
61:     for lineno, row in enumerate(csv.reader(f), start=1):
70:         cols = (0, 1)         # headerless: positional start,end

```

## 20. assistant (2026-06-22T09:33:49.269Z)

The engine claims hold up well. Key verifications:

- `cmd\_train` exists at `\_\_main\_\_.py:332` (doc says `:231` ? line number drifted, but the function exists).
- `synth\_trajectory\_from\_labels` at `stub\_runner.py:124`, reads `row["start\_time"]`/`row["end\_time"]` via `DictReader` at `:137-139` ? matches doc claim (doc said `:139`). ?
- `gate\_verdict` at `harness.py:142` (doc said `:98` ? drifted), `no\_paid\_llm\_training\_labels: True` at `:236` (doc said `:192`), `label\_source` "human (rally-annotator golden CSVs)" at `:222` (doc said `:178`). The line numbers have all drifted, but the facts hold.
- `gt\_loader` accepts `start`/`end` OR `start\_time`/`end\_time` at `:30` ? matches doc claim. ?

The doc's line-number citations are stale (engine has moved), but the doc itself flagged these are `[verified]` against code at a session in time. These are low-severity citation drift, not factual errors. Let me verify the gate\_verdict "refuses commercial ship" claim and the WASB native-runner "IN-FLIGHT / does not exist yet" claim, plus the annotator version.

## 21. user (2026-06-22T09:33:50.685Z)

```

142     def gate_verdict(eval_res: dict, floor: Optional[dict]) -> dict:
143         ""Pre-flight ship-gate.
144
145         Provisional F1 target (NOT owner-ratified; calibration tracked in #172).
146         Floor (necessary): heuristic LOVO ~0.480 (from quality table).
147         Multi-court ref (makes owned \"worth serving\" per data): 0.66.
148         Proposed for \"very high P/R\" + significance: F1@tIoU0.5 >= 0.70 on
149         multi-court folds with paired-sign p<0.05 vs heuristic (refine as corpus grows).
150         Beating floor + sign test necessary; C3 main commercial blocker.
151         ship=True only when all non-C3 blocks clear (honesty first).
152         ""
153         reasons: List[str] = []
154         floor_passed = None
155         sign = None
156         if floor:
157             floor_mean = float(floor["mean_f1"])
158             floor_passed = eval_res["mean_f1"] > floor_mean
159             if not floor_passed:
160                 reasons.append(f"LOVO mean {eval_res['mean_f1']:.3f} <= floor
{floor_mean:.3f}")
161             per = floor.get("per_video", {})
162             paired = [(r["f1"], per[r["video"]]) for r in eval_res["per_video"] if
r["video"] in per]
163             if paired:
164                 wins, n, p = sign_test([lv for lv, _ in paired], [fv for _, fv in paired])
165                 sign = {"wins": wins, "n": n, "p_one_sided": round(p, 4)}
166                 if p >= 0.05:
167                     reasons.append(f"paired sign test not significant ({wins}/{n} wins,
p={p:.3f}) "
168                                     f"? grow the corpus before claiming superiority")
169             else:
170                 reasons.append("no floor baseline provided ? nothing to beat")
171
172         # Provisional blocker (per owner review on #170; see #172 for calibration).
173         # (The number itself is not the blocker; calibration tracked in #172.)
174         reasons.append("ship-gate F1@tIoU target NOT YET CALIBRATED ? provisional; see issue

```

```

#172")
175     reasons.append("WASB checkpoint provenance C3 UNRESOLVED ? commercial ship blocked
regardless of F1")
176     # ship remains False until C3 resolved + other gates (by design).
177     return {"ship": False, "floor_passed": floor_passed, "sign_test": sign, "blockers":
reasons}
178
179
180     def run(manifest_path: str, out_dir: str, version: str,
181            floor_path: Optional[str] = None) -> dict:
182         with open(manifest_path, encoding="utf-8") as f:
183             specs = json.load(f)
184             floor = None
185             if floor_path:
186                 with open(floor_path, encoding="utf-8") as f:
187                     floor = json.load(f)
188
189             # 1) Honest LOVO eval (the number that matters).
190             eval_res = lovo_run(specs)
191
192             # 2) Final model on all videos.
193             clf, thr = train_final(specs)
194
195             # 3) Gate (pre-flight).
196             verdict = gate_verdict(eval_res, floor)
197
198             # 4) Artifact bundle.
199             bundle_dir = os.path.join(out_dir, version)
200             os.makedirs(bundle_dir, exist_ok=True)
201             weights_path = os.path.join(bundle_dir, "weights.json")

```

## 22. user (2026-06-22T09:33:55.064Z)

```

=== rally-annotator exists ===
 2 files changed, 11 insertions(+), 5 deletions(-)
1:--[[ rally_annotator.lua -- VLC Lua EXTENSION (v1.6.4)
14:  WHAT'S NEW IN v1.6.4
21:  WHAT'S NEW IN v1.6.3
vlc/rally_annotator.lua:12:
rally_number,start_time,end_time,ending_reason,sport,shots_count
vlc/rally_annotator.lua:105:    ..
"(rally_number,start_time,end_time,ending_reason,sport,shots_count; decimal seconds). "
vlc/rally_annotator.lua:128:local HEADER =
"rally_number,start_time,end_time,ending_reason,sport,shots_count\n"
docs/CSV_FORMAT.md:6:rally_number,start_time,end_time,ending_reason,sport,shots_count
docs/CSV_FORMAT.md:30:`rally_number,start_time,end_time` by header name and drops rows where
`end <= start`:
=== LICENSE ===
MIT License

```

## 23. assistant (2026-06-22T09:34:03.032Z)

Excellent ? all the engine and annotator core facts verify:

- `gate\_verdict` returns `{ "ship": False, ... }` with "WASB checkpoint provenance C3 UNRESOLVED ? commercial ship blocked regardless of F1" at `harness.py:175-177` ? matches doc's "ship-gate currently refuses a commercial ship regardless of F1; C3 unresolved + F1 not calibrated (#172)". ?
- Annotator `VERSION 1.6.4`, MIT, header  
`rally\_number,start\_time,end\_time,ending\_reason,sport,shots\_count` at `rally\_annotator.lua:128`  
? all match. ? (Note: rally-annotator repo is at `C:\Users\avidu\Projects\rally-annotator`, and  
`git pull` brought 11 insertions ? but version still 1.6.4.)

Now the key staleness question: the doc's §7 tracker. The doc is `DRAFT` status (line 3), but P0-P4 shipped (PRs #15, #22, #23) and screenshots shipped (#25) and the page was extensively localized (#74, #86, #87). The doc says "Last updated: 2026-06-20" but a lot landed after. Let

me verify the \$7 P-statuses precisely and check whether the doc has been superseded/should be IN PROGRESS or DONE. Let me check the nav/footer cross-links (P3) and confirm the "PR" column claims.

#### 24. user (2026-06-22T09:34:05.767Z)

```
=== nav cross-link to own-model in index.html ===
255:      <a href="/own-model" data-i18n="site.nav.ownModel">Own your model</a>
620:      <div><a href="/own-model">Own your model</a> · <a href="/capture">Capture
checklist</a> · <a href="mailto:hello@khelsutra.guru">hello@khelsutra.guru</a> · <a
href="https://khelsutra.guru">khelsutra.guru</a></div>
=== footer/nav in own-model.html (back link) ===
83:      <a class="brand" href="/">Khel<span class="play" aria-hidden="true"></span>Sutra</a>
98:      <div class="nav"><a href="/" data-i18n="site.ownModel.nav.home">? KhelSutra home</a><a
href="#annotate" data-i18n="site.ownModel.nav.annotate">Annotate</a><a href="#train"
data-i18n="site.ownModel.nav.train">Train your own model</a><a href="#setup"
data-i18n="site.ownModel.nav.setup">Setup & use</a></div>
191:      <span><a href="/">khelsutra.guru</a> · <a href="https://github.com/Khelsutra/rally-
annotator">rally-annotator (MIT)</a> · <span data-i18n="site.ownModel.foot.flyer">Academy
flyer:</span> <a href="/flyer" hreflang="en">English</a> · <a href="/flyer-kn" hreflang="kn"
lang="kn">?????</a> · <a href="/flyer-hi" hreflang="hi" lang="hi">?????</a></span>
=== capture.html nav has own-model link? ===
```

#### 25. user (2026-06-22T09:34:09.922Z)

```
=== PER_RALLY_SELECTION.md exists? line 122 ===
docs/PER_RALLY_SELECTION.md
A green SPA build + passing `parseQuery`/selection unit tests prove the **picker logic only** ?
**not** that the engine's rally boundaries are correct (the measured **rally-END** weakness is a
separate quality track, unaffected by this UI). "Conversational" here means a **deterministic
four-dimension parser**, not natural-language understanding. Nothing is implemented in `web/`
yet ? this is `[design]`. The feature curates *what the engine already found*; it cannot surface
a rally the detector missed, nor split/merge a mis-bounded one (that's the M4 correction loop).
```

### 7. Deliverables & progress tracker ? \*\*source of truth\*\*

Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. \*\*One small PR per row\*\*, branched from `origin/master`, no stacking. CI (`npm test`) must stay green; merge to `master` auto-deploys the site (docs/`web/` changes don't touch the live `site/` Worker).

```
=== LOCAL_APPLIANCE_ARCHITECTURE.md exists? (doc says lands with PR #13) ===
docs/LOCAL_APPLIANCE_ARCHITECTURE.md
=== web SPA router? ===
=== web/src/api/client.ts exists? ===
web/src/api/client.ts
web/src/api/types.ts
web/src/lib/log.ts
```

#### 26. assistant (2026-06-22T09:34:15.511Z)

Good. Cross-links confirmed (P3 ? accurate ? nav line 255, footer line 620 in index.html). `LOCAL\_APPLIANCE\_ARCHITECTURE.md` exists now (the doc said it "lands with PR #13" ? it has landed, so that forward-looking note is stale but harmless). SPA helper files exist. Let me check the SPA router claim ? doc says "today web/src/App.tsx is one static screen, no router". Let me verify, plus check whether `PATCH /api/segments` endpoint now exists in the engine (P8 claimed unbuilt).

#### 27. user (2026-06-22T09:34:18.012Z)

```
=== react-router in web/package.json? ===
NO router dep
=== App.tsx routing? ===
8:import { LanguageSwitcher } from "../components/LanguageSwitcher";
11:import { SetupWizard, isWizardDone } from "../components/SetupWizard";
64: // P2c: the first-run wizard shows once per browser, before the ingest panel, when
```

```

nothing's loaded.
65:  const [wizardDone, setWizardDone] = useState(isWizardDone);
227:      <LanguageSwitcher />
240:      {/* P2c first-run wizard: before anything is loaded, walk a new user through AI-key
/ compute /
242:      {!videoId && !wizardDone && <SetupWizard onClose={() => setWizardDone(true)} />} />}
246:      {!videoId && wizardDone && (

```

## 28. user (2026-06-22T09:34:19.852Z)

```

=== PATCH /api/segments endpoint exists? ===
backend/eval/experiment.py:18: `candidate_segments.stats`) from the **cheap, swappable DECISION**
(given those
backend/eval/experiment.py:36: `Database.query_candidate_segments`.
backend/eval/experiment.py:628:  """Bridge persisted detection ? `VideoCandidates` via
`Database.query_candidate_segments`.
backend/eval/experiment.py:635:  rows = db.query_candidate_segments(video_id,
detector=detector)
backend/eval/harness.py:7:  - ``candidates`` ? raw detector windows (`candidate_segments`
table). What
backend/eval/harness.py:58:  rows = db.query_candidate_segments(video_id,
detector=detector)
backend/eval/training.py:26:# YOLO Stage-1 persists these per candidate window in
`candidate_segments.stats`
backend/eval/training.py:34:  """Feature vector for one YOLO candidate row (dict from
query_candidate_segments).
backend/infrastructure/database.py:217:  CREATE TABLE IF NOT EXISTS candidate_segments
(
backend/infrastructure/database.py:238:  ON candidate_segments(video_id, detector)
=== own-model engine route check done ===

```

## 29. assistant (2026-06-22T09:34:29.974Z)

Findings on these:

- \*\*`react-router-dom` is still NOT a dep\*\* in `web/package.json` ? matches doc claim "no router". ? But the doc says "today `web/src/App.tsx` is one static screen" ? it is now MORE than one static screen: it has a `SetupWizard` (P2c, PR #88) gated on `!videoId && !wizardDone`. So the "one static screen" framing is mildly stale (the SPA grew a wizard), but the core claim (no router, no correction screen) still holds ? low severity.
- \*\*`PATCH /api/segments/{id}` endpoint still does not exist\*\* in the engine ? confirmed via grep (only `candidate\_segments` table refs, no PATCH route, `human\_label` columns still dead). Doc P8 (? Gated) is accurate. ?

Now the central staleness assessment. The doc header says `Status: DRAFT` and `Last updated: 2026-06-20`, but a substantial amount shipped after that date (the page #22, screenshots #25, localization #74/#86/#87, org-link change #23). Let me confirm the \$7 P5 status is genuinely stale (screenshots done but marked ?) and check whether the doc was ever updated past 2026-06-20.

## 30. user (2026-06-22T09:34:33.003Z)

```

=== full git history of the doc ===
87520b8 chore(site): point rally-annotator links at the new Khelsutra org (#23)
b1110e1 feat(site): own-model page ? annotate your footage & train your own model (annotate-
train P1-P4) (#22)
6b62ebf docs(annotate-train): plan rally-annotator integration + "train your own model" USP
pages (#15)
=== last commit date that touched it ===
2026-06-20 20:24:26 +0530 chore(site): point rally-annotator links at the new Khelsutra org
(#23)

```

## 31. user (2026-06-22T09:34:34.996Z)

```

=== is the doc referenced anywhere (NEXT_STEPS, current-state)? ===

```

```
docs/ANNOTATE_AND_TRAIN.md
docs/archives/FRIENDLY_AND_MULTILINGUAL.md
=== NEXT_STEPS own-model/annotate mentions ===
31: footage + a dev "P5" badge ? clean up before public embed), and/or produce a polished
Grok/Gemini cut.
36: `parseQuery()` query bar (#28), P4 SelectionFooter + `min_shot_count` warning (#30), P5
real load +
58: retention, cost ceiling, GPU host, P5 bridge).
```

### 32. assistant (2026-06-22T09:34:44.894Z)

Now I have the full picture. The doc was last touched in PR #23 (2026-06-20), and its §7 tracker reflects state as of that PR (P1-P4 ?, P5 ?). But since then:

- **PR #25 shipped real screenshots** ("show the real annotator screenshots") ? `annotate-enable.png`, `annotate-tool.png` (133KB/153KB real PNGs). So §7 **P5** ("Capture + commit real screenshots... replace placeholders", marked ? Owner-gated) is **stale** ? the annotation-tool screenshots are now real and committed. (Note: P5 also mentions footage screenshots with people, which remain placeholder/deferred ? but the "replace placeholders" with real annotator UI shots is done.)
- The page got heavily extended after the doc froze: reel gating (#24/#25), org-link migration (#23 ? captured in doc), academy flyer link (#44), and **full 6-language localization** (#70, #74, #86, #87) which the doc never anticipated.
- The `docs/README.md` index does NOT list `ANNOTATE_AND_TRAIN.md` ? so the doc's repeated tracking claim ("indexed in `docs/README.md`", DoD: "doc indexed in `docs/README.md`", §7 P0: "indexed in docs/README.md") is unverified/false.
- The doc is still `DRAFT` though P0-P4 are shipped/merged ? at minimum it should be `IN PROGRESS`.

Let me verify whether the doc is indexed in docs/README.md by reading the whole file (I only saw the first 40 lines), and confirm the `og.png`/`make_og.py`` reuse claim.

### 33. user (2026-06-22T09:34:47.517Z)

```
=== full docs/README.md grep for annotate/own/train ===
NOT FOUND in docs/README.md
=== og.png + make_og.py exist? ===
site/make_og.py
site/public/og.png
```

### 34. assistant (2026-06-22T09:34:55.796Z)

Confirmed: `ANNOTATE_AND_TRAIN.md` is NOT in `docs/README.md`. And `og.png`/`make_og.py`` exist (the reuse-map precedent is accurate ? though the doc cited `site/og.png``, it's actually `site/public/og.png``; minor path drift, `site/make_og.py`` is correct).

I have enough evidence. Let me compile the findings. Summary of the staleness picture:

- The doc is a **DRAFT** tracker whose marketing-page work (P0-P4 + P5 screenshots) is essentially shipped and merged, but the tracker has not been updated since 2026-06-20 (PR #23) and is now lagging reality (P5 screenshots done; status still DRAFT; not indexed; line-number citations drifted; SPA grew a wizard).
- Core engine/annotator technical claims (schema bridge, ship-gate, MIT annotator, guardrails) all still **verify as correct** ? only line numbers drifted.

This is a tracker that is mostly complete but not closed out ? minor-to-significant staleness, an archive candidate once honestly updated.

### 35. user (2026-06-22T09:35:43.673Z)

Structured output provided successfully

### 36. assistant (2026-06-22T09:35:56.569Z)

Audit complete. I verified every concrete claim in `docs/ANNOTATE_AND_TRAIN.md` against the live KhelSutra repo (master @ #94), the engine repo (`badminton-highlight-indexer``), and the `rally-`



annotator` repo.

Verdict: **\*\*significant-staleness\*\*** ? a DRAFT tracker whose work is largely shipped but never closed out.

High-severity (a reader would act on these):

- **\*\*P5 screenshots marked Todo but actually shipped\*\*** ? PR #25 committed real annotator-UI screenshots (`site/public/img/annotate-enable.png`, `annotate-tool.png`), referenced in `own-model.html:111-112`.
- **\*\*"Indexed in docs/README.md" is false\*\*** ? claimed three times (header, §7 P0, DoD), but `ANNOTATE\_AND\_TRAIN.md` does not appear anywhere in `docs/README.md`.

Medium:

- Header still says `DRAFT` / `Last updated: 2026-06-20` though P0-P4 are merged (should be IN PROGRESS, or DONE+archive).
- §7 PR column uses `(this PR)` placeholders instead of the real merged PRs (#22/#23).

Low (citation drift / forward-looking notes, facts still correct):

- Every engine `file:line` citation has drifted (`gate\_verdict` :98?:142, `cmd\_train` :231?:332, attestation :192?:236, label\_source :178?:222) ? the underlying facts all still verify.
- "one static screen, no router" is now stale (App.tsx grew a `SetupWizard`, PR #88) ? though "no router" is still true.
- `LOCAL\_APPLIANCE\_ARCHITECTURE.md` "lands with PR #13" ? it already exists.
- OG asset path is `site/public/og.png`, not `site/og.png`.

The core technical substance (schema bridge, ship-gate refusing a commercial ship, MIT annotator v1.6.4, guardrails, unbuilt `PATCH /api/segments`) all verified **\*\*correct\*\***. Once honestly updated, this is an archive candidate.