

claude-windows / 8d4866bf-8edd-4da9-909e-022ddee12675

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-inde  
xer\8d4866bf-8edd-4da9-909e-022ddee12675\subagents\workflows\wf_f9255361-69d\agent-  
a28b2d0e1d2283cea.jsonl`  
- SHA-256: `43ef6b2dedbcd14998c0706aa7b4ab88e8961a642474aedc2a56e83c004637f3`  
- Source modified: `2026-06-22T09:42:58+00:00`  
- Imported at: `2026-07-05T16:48:25+00:00`  
- project: `wf_f9255361-69d`  
- session_id: `8d4866bf-8edd-4da9-909e-022ddee12675`
```

Transcript

1. user (2026-06-22T09:32:46.021Z)

You are auditing ONE doc in the KhelSutra brand/site repo for staleness, as part of a "are all the docs up to date?" sweep.

Repo root: C:\Users\avidu\Projects\kheelsutra (bash path: /c/Users/avidu/Projects/kheelsutra)
The repo is on `master`, freshly synced to remote HEAD ? the latest merged PR is #94. Today is 2026-06-22.

DOC TO AUDIT: deploy/DEPLOY_ALPHA.md

Your job: read the doc IN FULL, extract every concrete, falsifiable claim, and VERIFY each against the LIVE repo ? do NOT trust the doc's own wording. Hard discipline rule: a claim is "correct" ONLY when you have file:line or command-output evidence for it; if you cannot verify it, mark verdict "unverifiable" ? never guess, never assume the doc is right. Re-check against actual code/git.

Kinds of claims to check:

- File / path / route references ("served at /capture", "see src/store.js", "public/og.png", "systemd kheelsutra-site") -> does the path or route ACTUALLY exist now? Confirm with git ls-files / Read / the routing in site/server.js + site/src/index.js.
- PR numbers and counts ("PRs #1-#8 merged", "20 tests green", "6 languages", "3-step wizard", "P0..P5 done") -> verify against `git -C <root> log`, the actual test output/count, and the code.
- Status markers (DONE / pending / IN PROGRESS / DRAFT / "?" / "owner-gated" / "NEXT") -> still accurate vs the code & git history? Is something marked pending that is actually shipped, or vice-versa?
- Dates, "live at <url>", "merged in #N".
- Whether the doc describes something now SUPERSEDED or fully shipped -> is it an archive candidate (a tracker whose work is complete, or content duplicated by docs/archives/)? Note: docs/archives/FRIENDLY_AND_MULTILINGUAL.md already exists as the archived multilingual tracker.

Tools: Read the doc + every cited file; Grep/Glob to locate code; Bash for `git -C /c/Users/avidu/Projects/kheelsutra log --oneline`, `... ls-files`, and greps. To check a test count, you may run the suite (`cd /c/Users/avidu/Projects/kheelsutra/site && node --test` or `npm test`), but prefer cheap checks. Verify against master as it stands NOW.

Severity: high = a factually wrong claim a reader would act on (wrong path/route, a "DONE" that isn't, a wrong count/PR that misleads). medium = outdated-but-harmless (stale PR range, old "where we are", superseded framing). low = cosmetic.

Do NOT modify any file ? audit only. Return ONLY the structured object (doc, overallStatus, findings[], summary).

2. assistant (2026-06-22T09:32:49.294Z)

I'll audit the deploy/DEPLOY_ALPHA.md doc. Let me start by reading it in full and gathering repo context.

3. user (2026-06-22T09:32:49.691Z)

```
1      # Deploy KhelSutra Studio for alpha ? a shareable page, no SSH
2
3      > **Status:** `READY` (owner executes the cloud steps) . **Topology:** split ? SPA on
Cloudflare Pages + engine on an `api.*` cloudflared tunnel, Cloudflare Access over both (the
engine's `HOSTING_PLAN.md` $2 alpha plan). . **Cost:** ~$0/mo standing (Pages + Tunnel + Access
are free ?50 users) + Gemini per-run (~$0.05/match).
4
5      This turns the local-only flow (run engine + SPA on your box, drive it yourself) into a
**page you can share with alpha testers** ? they open `app.khelsutra.guru`, sign in (email
allowlist), point at a video, and get a reel. No SSH, no home IP exposed, no inbound ports.
6
7      ## 0. The shape
8
9      ```
10     tester browser
11     ? https
12     ?
13     Cloudflare edge (free): DNS . TLS . Access (email-OTP allowlist) in front of BOTH
hostnames
14     ??? app.khelsutra.guru ? Cloudflare Pages          (the web/ SPA; built with
VITE_API_BASE)
15     ??? api.khelsutra.guru ? cloudflared Tunnel (outbound-only) ?? 127.0.0.1:8000 on
YOUR box
16
17
18
19
20
21     - **No inbound ports.** `cloudflared` dials out; the engine stays bound to
`127.0.0.1:8000`. The box's home/public IP is never exposed.
22     - **Auth at the edge.** Cloudflare Access gates both hostnames; the engine *also*
verifies the `Cf-Access` JWT in-app (defence in depth ? `security.edge_auth_enabled`), so a
direct hit to the tunnel can't spoof identity.
23     - **Cross-origin, on purpose.** `app.` (SPA) and `api.` (engine) are different origins ?
the SPA is built with `VITE_API_BASE=https://api.khelsutra.guru/api` (PR #64) and the engine
locks CORS to the SPA origin via `RALLY_CORS_ALLOW_ORIGINS` (no engine code change ? it's
already env-driven).
24
25     > **Alternative (single-origin):** if you'd rather run **one box** that serves the SPA
*and* the API behind one origin (Caddy reverse-proxy, no Pages, no cross-origin), see
`LOCAL_APPLIANCE_ARCHITECTURE.md` D1. This runbook ships the **split** topology you chose
(stable Pages-hosted page + a thin API tunnel).
26
27     ---
28
29     ## 1. Prerequisites
30
31     - A **box** to run the engine: the CPU droplet `168.144.126.176` (Gemini-only ? no GPU;
fine for alpha) **or** your Windows+WSL GPU box (real WASB). `ffmpeg` + Python 3.8+ installed;
the `badminton-highlight-indexer` engine checked out.
32     - A **Cloudflare account** with the `khelsutra.guru` zone (you already run the marketing
site there).
33     - `cloudflared` installed on the box (`https://developers.cloudflare.com/cloudflare-
one/connections/connect-networks/downloads/`).
34     - A **fresh** `GEMINI_API_KEY` (rotate the chat-shared one ? it uploads downsized chunks
to Google; this is a paid + cloud action).
35
36     ---
```

```

37
38     ## 2. Engine (the `api.` side)
39
40     ### 2.1 Install (with the auth extra)
41     ```bash
42     cd badminton-highlight-indexer
43     pip install -e .[auth]          # PyJWT[crypto] ? REQUIRED when edge auth is on, else
every authenticated request is rejected 401
44     ```
45
46     ### 2.2 The safe-exposed config
47     Create / edit `config.json` on the box. **Where it must live:** the engine reads
`$RALLY_APP_HOME/config.json` if `RALLY_APP_HOME` is set, else `./config.json` relative to *the
directory you launch the server from* ? **not** necessarily the repo root. If the engine reads a
different (or absent) file it silently falls back to all-defaults, which means **edge auth is
OFF**. To make writer and reader agree by construction, write it with `python -m backend.main
--setup` (it targets the exact path the server reads) and edit *that* file:
48     ```jsonc
49     {
50         "security": {
51             "edge_auth_enabled": true,                // turn the CF-Access JWT verify ON
52             "allowed_video_root": "/srv/khelsutra/incoming", // the path-jail ? REQUIRED when
exposed
53             "upload_dir": "/srv/khelsutra/incoming/uploads", // MUST sit UNDER
allowed_video_root (see §7)
54             "cf_team_domain": "https://<your-team>.cloudflareaccess.com",
55             "cf_access_aud": "<the CF Access application AUD>" // filled in §4
56         },
57         "indexing": {
58             "default_segmenter": "gemini" // ? CPU box MUST set this ? the shipped default
`yolo_hybrid` needs torch/GPU, and the SPA sends no segmenter, so every job fails "No module
named 'torch'" without it
59         }
60     }
61     ```
62     And export the env (a `.env` or your shell):
63     ```bash
64     export GEMINI_API_KEY="<your-rotated-key>"
65     export RALLY_CORS_ALLOW_ORIGINS="https://app.khelsutra.guru" # lock CORS to the SPA
origin
66     export RALLY_ALLOWED_HOSTS="api.khelsutra.guru" # REQUIRED ? see ? below;
without it EVERY tunneled request 400s
67     export RALLY_BIND_HOST="127.0.0.1" # keep the loopback bind
(cloudflared dials 127.0.0.1:8000)
68     ```
69
70     > **? Two env vars you MUST set, or the tunnel is dead-on-arrival.** With
`edge_auth_enabled: true`, the engine (a) flips its bind to `0.0.0.0` and (b) keeps the DNS-
rebinding **Host-header guard** at its loopback-only default ? so it boots "healthy" but then
**400's every request `cloudflared` forwards with `Host: api.khelsutra.guru` (including the
§6 `/api/health` check). The boot only *warns* about this (`EXPOSED_HOST_ALLOWLIST_LOOPBACK`),
it does not stop. Set **`RALLY_ALLOWED_HOSTS="api.khelsutra.guru"` (the load-bearing fix ? the
Host allowlist is independent of the bind) and **`RALLY_BIND_HOST="127.0.0.1"` (so the bind
stays loopback as §0/§6 claim; `cloudflared` dials loopback anyway). *Verified against engine
`origin/master`: with the allowlist unset, `GET /api/health` with `Host: api.khelsutra.guru` ?
`400`; with it set ? `200`.
71
72     > **The engine refuses to start unsafe.** The boot posture (engine PR #294) **fails
fast** if `edge_auth_enabled` is on but the `allowed_video_root` jail or
`cf_team_domain`/`cf_access_aud` is missing, and **warns** if `PyJWT` isn't installed. So a
half-configured exposure can't silently come up ? fix what it prints and restart.
73
74     ### 2.3 Run
75     ```bash
76     python -m backend.main --server # binds 127.0.0.1:8000 ? kept loopback by

```

```

RALLY_BIND_HOST ($2.2); with edge_auth ON and that var unset it would bind 0.0.0.0 instead
77     ``
78     Keep it alive with a process manager (`systemd` on the droplet, mirroring
`sites/scripts/provision-vm.sh`'s unit; or `pm2`/`nssm` on Windows).
79
80     ---
81
82     ## 3. The SPA on Cloudflare Pages (the `app.` side)
83
84     The SPA ships as a static Pages site, git-connected like the marketing site (auto-
deploy on merge to `master`).
85
86     1. Cloudflare dashboard ? Workers & Pages ? Create ? Pages ? Connect to Git ? pick
the `khelsutra` repo.
87     2. Build settings:
88         - Root directory: `web`
89         - Build command: `npm run build`
90         - Build output directory: `web/dist`
91         - Environment variable: `VITE_API_BASE = https://api.khelsutra.guru/api` ? this
is what points the SPA at the engine tunnel (PR #64).
92     3. Custom domain: add `app.khelsutra.guru` to the Pages project.
93
94     > Every push to `master` now rebuilds + redeploys the SPA. To preview a change first,
Pages builds a per-PR preview URL.
95
96     ---
97
98     ## 4. Cloudflare Access (gate BOTH hostnames)
99
100    1. Zero Trust ? Access ? Applications ? Add ? Self-hosted.
101    2. Application domain: add both `app.khelsutra.guru` and
`api.khelsutra.guru` to the same application (one app ? one shared session cookie, so a tester
signs in once and both the SPA and its API calls are authorized).
102    3. Policy: Action Allow, rule Emails ? your alpha testers' addresses (free ? 50
users ? fits 10?25 testers). Use one-time-PIN (email OTP) so testers need no Cloudflare account.
103    4. Copy the application `AUD` tag ? paste into the engine's `config.json`
`security.cf_access_aud` ($2.2), and set `cf_team_domain` to `https://<your-
team>.cloudflareaccess.com`. Restart the engine.
104    5. CORS preflight ?: the browser sends an unauthenticated `OPTIONS` preflight to
`api.` for the cross-origin POSTs. In the Access application's Settings ? enable ?Bypass
options requests to CORS preflight? (or add an explicit OPTIONS bypass), else the preflight is
redirected to the login page and every API call fails. The engine's CORS middleware
(`allow_credentials=false`, no cookies) then answers the preflight.
105
106    ---
107
108    ## 5. The tunnel (engine ? Cloudflare)
109
110    ```bash
111    cloudflared tunnel login                # browser auth to your zone
112    cloudflared tunnel create khelsutra-studio # note the <TUNNEL_UUID>
113    # copy deploy/cloudflared.config.example.yml ? /etc/cloudflared/config.yml, fill
<TUNNEL_UUID>
114    cloudflared tunnel route dns khelsutra-studio api.khelsutra.guru
115    cloudflared tunnel run khelsutra-studio # or install as a service:
`cloudflared service install`
116    ```
117    See `deploy/cloudflared.config.example.yml` for the ingress (only `api.khelsutra.guru` ?
http://127.0.0.1:8000).
118
119    ---
120
121    ## 6. Verify (the posture test + the CUJ)
122
123    Exposure posture (must hold ? this is what makes it safe to share):

```

```

124 - From OUTSIDE the box: `curl --max-time 5 http://<box-public-ip>:8000/api/health` ?
**refused / timeout** (the engine is `127.0.0.1`-bound; no inbound port). ? The box is not
directly reachable.
125 - `curl -i https://api.khelsutra.guru/api/health` **without** a CF session ? a **302 to
the Access login** (not a 200). ? The gate is in front.
126 - Open `https://app.khelsutra.guru` in a browser ? CF Access login ? after sign-in, the
SPA loads and `/api/health` returns 200 (same session cookie). ?
127
128 **The shareable CUJ (what a tester does):** open `app.khelsutra.guru` ? sign in ?
**upload a video from their device** (the SPA chunked-upload lands it in `upload_dir`, inside
the jail) ? **Process** ? watch progress ? pick rallies ? **Build** ? **download the reel**. All
from the page. *(Pasting a `gs://` URI does NOT work on this box ? see §7: the
`allowed_video_root` jail rejects every non-local URI. Browser upload is the alpha path.)*
129
130 ---
131
132 ## 7. Honest caveats / deferred
133
134 - **`upload_dir` must sit UNDER `allowed_video_root`.** If you enable browser upload, an
upload landing outside the jail is rejected by `/api/process` (a functional break, not a hole).
Keep `upload_dir` inside the jail (§2.2).
135 - **`gs://` / `s3://` / `gdrive://` URIs are REJECTED on the exposed box.** Setting
`allowed_video_root` (required to expose, §2.2) path-jails `/api/process` to *local* paths under
that root; the engine resolves a `gs://` URI to a local path that escapes the jail ? `400`
(pinned by engine `tests/test_api_endpoints.py::test_process_hosted_mode_rejects_nonlocal_uri`).
Cloud-URI ingest works **only** when `allowed_video_root` is unset (the local single-user
posture ? the opposite of this runbook). For alpha, the ingest path is **browser upload** (it
lands inside the jail). *(Scheme-exempting recognized buckets from the jail would be an engine
change, not a config flip.)*
136 - **You MUST set `edge_auth_enabled=true` to expose.** The boot posture (and the in-app
JWT verify) only engage then. Do not put the engine behind the tunnel with edge auth off.
137 - **Cross-origin reel download:** the reel streams from `api.` with `Content-
Disposition: attachment`, so it downloads even cross-origin ? but verify the filename in your
browser during the CUJ (cross-origin `` filename hints are ignored; the header is
the authority).
138 - **Upload size:** Cloudflare's free proxy caps request bodies at **100 MB** ? large
matches upload as ~95 MB chunks (the SPA chunked-upload path honours the engine's `max_part_mb`,
default **95** ; lower `security.max_part_mb` to shrink them). `gs://` bucket pre-staging is
**not** an option on this box (see the jail caveat above).
139 - **Box offline = engine offline.** The Pages SPA stays up and shows a friendly "engine
offline" banner (`/api/health` ping); processing resumes when the box is back. Beta removes this
with a queue (`HOSTING_PLAN.md` §3).
140 - **Gemini = paid + uploads footage to Google.** Keep the key rotated; use it only for
clips you're OK sending.
141
142 ---
143
144 ## 8. Teardown
145 `cloudflared tunnel cleanup khelsutra-studio` + delete the tunnel; remove the Pages
custom domain; disable the Access application; stop the engine service. The box's `output/` +
`uploads/` retain originals ? sweep them per your retention policy before any wider exposure.
146
147 ---
148 ### Changelog
149 - 2026-06-22 ? **correctness pass** (adversarial verify against engine `origin/master` +
a hermetic repro). Fixed two would-break-the-deploy gaps: (1) **`RALLY_ALLOWED_HOSTS` was
missing** ? with `edge_auth` on, the Host-header guard 400s every tunneled request until it's
set (added it + `RALLY_BIND_HOST` to §2.2; repro confirmed 400?200); (2) **`gs://` ingest is
rejected** by the `allowed_video_root` jail on the exposed box (dropped the `gs://` promise from
§6/§7; alpha path = browser upload). Minor: 401-not-500 wording (§2.1), `config.json`
location/`--setup` note (§2.2), bind comment (§2.3), ~95 MB chunk size (§7).
150 - 2026-06-21 ? created. Split-topology alpha deploy kit (cloudflared + Pages + CF
Access) consuming the engine exposure-safety boot posture (engine #294) + the SPA
`VITE_API_BASE` (#64). Pairs with `LOCAL_APPLIANCE_ARCHITECTURE.md` (single-origin alternative)
+ the engine `HOSTING_PLAN.md` §2.

```

4. user (2026-06-22T09:32:50.359Z)

50b2bd7 chore(repo): ignore local runtime artifacts + correct stale NEXT_STEPS claims (#94)
 e2570e4 docs: archive the Friendly & Multilingual tracker (engineering complete) (#93)
 0bef61e docs(friendly-multilingual): close out the chapter ? log #89/#90/#91 + strengthen fitNav tests (#92)
 9f77ddc test(site): browser layout-regression CI guard + auto-collapse the nav to fit every locale (#91)
 4415fe8 feat(site): localize remaining home sections (6 langs) + smooth the responsive header (#90)
 26c9da1 feat(site): localize the hero demo mockup + #how feature block (6 languages) (#89)
 7ae9cee feat(web): P2c first-run setup wizard (3 steps from /api/capabilities, 6 languages) (#88)
 cf59fbf feat(site): add es/da/id to the marketing site (parity with the SPA's 6 languages) (#87)
 17242a9 feat(site): localize own-model body content to kn/hi (#86)
 7b04f02 deploy(kit): align with the live-verified droplet deploy (+ the gemini-default fix) (#85)
 5e5a2ac docs(friendly-multilingual): close out the chapter (\$9 DoD + UI/UX + B7 deferred) (#84)
 4c0f35c fix(site): UI/UX ? compact language switcher + header overflow fix + localized hero (#83)
 0fae364 feat(web+site): A2 ? localized WhatsApp share for the reel (SPA + site) (#82)
 1e52cfd deploy(droplet): tailored alpha kit for the CPU droplet (engine api. side) (#81)
 64076a5 test(web): add a coverage floor (CI-gated) + i18n regression tests for the #77 audit (#80)
 8328721 docs(tracker): refresh the \$7 milestone prose to merged reality (A6/A7/#77 landed) (#79)
 3f0cdb8 docs(deploy): fix two would-break-the-deploy gaps in the alpha runbook (RALLY_ALLOWED_HOSTS + gs:// jail) (#78)
 9f1ec9c fix(i18n): native-review pass on the kn/hi SPA console catalogs (#77)
 9a2d7d2 feat(site): A7 ? camera-placement diagram on the capture page (#76)
 aa8d5e9 feat(site): A6 ? plain-language Request-Access GPU/storage questions (#75)
 1cd84af feat(site): A3 ? localize recording-education + marketing to kn/hi, Noto on the site (#74)
 c1848df feat(web): LW2 ? send the UI locale to /api/compile (watermark) + /api/process (analytics) (#73)
 1b1fa7b feat(web+site): B3 ? trilingual marketing site via shared t()/data-i18n shim (D-CATALOG-HOME) (#70)
 f7e2bd5 chore(site): add robots.txt + a "crawling / bot protection" runbook for khelsutra.guru (#71)
 bea8a2a docs(appliance): flip T11 retention sweep to ? ? tool shipped (engine #299), scheduling tracked (#301) (#69)
 46e0244 feat(web): A9 ? localized tap-only filter controls (#68)
 a2292ef feat(web): MULTILINGUAL Phase A ? Spanish/Danish/Indonesian SPA catalogs (#67)
 ca619fb feat(web): T4 ? in-UI file upload (the runnable-tool keystone) (#66)
 4627830 docs(deploy): alpha deploy kit ? make Studio a shareable page (split: Pages + api.* tunnel + CF Access) (#65)
 8bbc557 feat(web): VITE_API_BASE ? let the SPA target a split api.* origin (default same-origin /api) (#64)
 92a6c45 feat(web): A8 ? background-processing reassurance + icon+text status chips (#63)
 501c014 feat(web): A5 ? friendly engine-error mapping across build + load (shared errors.*) (#62)
 95195a5 feat(web): A4 ? mobile touch targets (?44px) + visible tooltip text (#61)
 ccdcfcd feat(web): A1 ? plain-language copy pass on the en catalog (de-jargon) (#60)
 a20002f docs: propose Multilingual Expansion ? 7 new locales (th/id/da/es/zh-Hans/ja/ko) (#58)
 a0475bc docs: hygiene refresh ? sync trackers to shipped reality (#59)
 7750afc docs(dry-run): Path-B live e2e DONE ? in-UI bucket?reel CUJ + screencast (DoD \$6 box 3 ?) (#57)
 4525c5d feat(web): B-P2 ? in-UI bucket ingest (point at a gs:// video ? process ? load) (#56)
 d0461fd docs: correct stale "bucket ingest is net-new" claims (B-P1 already shipped in M2/#280) (#55)
 2421557 docs(dry-run): record Path-A local result (en+hi reel downloaded) + the engine fix it surfaced (#54)

5. assistant (2026-06-22T09:32:53.816Z)

The doc references many engine-side claims (badminton-highlight-indexer) and site/web repo claims. Let me verify the falsifiable claims that live in THIS repo (khelsutra). Let me check the cited files, PR numbers, and routes.

6. user (2026-06-22T09:32:56.601Z)

```
=== deploy/ files ===
deploy/DEPLOY_ALPHA.md
deploy/cloudflared.config.example.yml
deploy/droplet/README.md
deploy/droplet/config.json.example
deploy/droplet/engine.env.example
deploy/droplet/khelsutra-engine.service
=== web/ structure ===
web/.gitignore
web/README.md
web/index.html
web/package-lock.json
web/package.json
web/src/App.tsx
web/src/api/client.test.ts
web/src/api/client.ts
web/src/api/locale.test.ts
web/src/api/types.ts
web/src/api/upload.test.ts
web/src/components/CompileResult.tsx
web/src/components/IngestPanel.test.tsx
web/src/components/IngestPanel.tsx
web/src/components/LanguageSwitcher.tsx
web/src/components/QueryBar.examples.test.ts
web/src/components/QueryBar.test.tsx
web/src/components/QueryBar.tsx
web/src/components/RallyGrid.tsx
web/src/components/SelectionFooter.test.tsx
web/src/components/SelectionFooter.tsx
web/src/components/SetupWizard.test.tsx
web/src/components/SetupWizard.tsx
web/src/components/StatusChip.tsx
web/src/components/TapFilters.test.tsx
web/src/components/TapFilters.tsx
web/src/components/a4.touch.test.tsx
web/src/components/a8.status.test.tsx
web/src/components/build-error.test.tsx
web/src/components/upload.test.tsx
=== LOCAL_APPLIANCE_ARCHITECTURE ===
deploy/cloudflared.config.example.yml
docs/LOCAL_APPLIANCE_ARCHITECTURE.md
```

7. user (2026-06-22T09:32:58.580Z)

```
=== PR #64 ===
8bbc557 feat(web): VITE_API_BASE ? let the SPA target a split api.* origin (default same-origin /api) (#64)
=== PR #65 ===
4627830 docs(deploy): alpha deploy kit ? make Studio a shareable page (split: Pages + api.* tunnel + CF Access) (#65)
=== PR #78 ===
3f0cdb8 docs(deploy): fix two would-break-the-deploy gaps in the alpha runbook (RALLY_ALLOWED_HOSTS + gs:// jail) (#78)
=== PR #294 (engine) ===
```

8. assistant (2026-06-22T09:33:03.661Z)

PR #64, #65, #78 confirmed. PR #294 is an engine PR (not in this repo). Let me verify the cited

files: `deploy/cloudflared.config.example.yml`, `docs/LOCAL\_APPLIANCE\_ARCHITECTURE.md` (D1 reference), `HOSTING\_PLAN.md` (which is NOT in this repo ? it's the engine's), the `site/scripts/provision-vm.sh`, and the `VITE\_API\_BASE` claim.

9. user (2026-06-22T09:33:04.567Z)

```
1      # cloudflared tunnel config ? KhelSutra Studio ENGINE (api.khelsutra.guru ? FastAPI on
the box).
2      #
3      # Copy to the engine host (e.g. /etc/cloudflared/config.yml or
~/.cloudflared/config.yml), fill in
4      # <TUNNEL_UUID>, then run cloudflared as a service. The tunnel dials OUT to Cloudflare,
so NO inbound
5      # port is opened on the box ? the engine stays bound to 127.0.0.1:8000 (its safe
default), and the
6      # only way in is through Cloudflare Access. See DEPLOY_ALPHA.md.
7      #
8      # The SPA (app.khelsutra.guru) is NOT served here ? it ships on Cloudflare Pages
(DEPLOY_ALPHA.md §3).
9      # This tunnel carries ONLY the engine API.
10
11     tunnel: <TUNNEL_UUID>
12     credentials-file: /etc/cloudflared/<TUNNEL_UUID>.json
13
14     ingress:
15         # The engine FastAPI control plane (POST /api/process, GET /api/jobs, /api/query,
/api/compile,
16         # GET /api/highlights, /api/health). Bound to 127.0.0.1:8000 on the box ? never
0.0.0.0.
17         - hostname: api.khelsutra.guru
18           service: http://127.0.0.1:8000
19           originRequest:
20               # The async pipeline (fetch ? segment ? stitch) can run for minutes; don't let the
tunnel
21               # time the request out before the job is submitted (submit returns 202 fast, but
be generous).
22               connectTimeout: 30s
23               noTLSVerify: false
24
25         # A cloudflared ingress MUST end with a catch-all.
26         - service: http_status:404
27
```

10. user (2026-06-22T09:33:07.013Z)

```
=== provision-vm.sh ===
site/scripts/provision-vm.sh
=== site/scripts ===
site/scripts/copy-locales.mjs
site/scripts/droplet-smoke.sh
site/scripts/parity-check.sh
site/scripts/provision-vm.sh
site/scripts/teardown-vm.sh
=== HOSTING_PLAN ===
(none above = not in this repo)
=== VITE_API_BASE usage ===
web/src/api/client.test.ts:27: it("defaults to same-origin /api when VITE_API_BASE is unset",
() => {
web/src/api/client.test.ts:31: it("uses VITE_API_BASE for a split deploy (SPA on Pages + api.*
tunnel)", () => {
web/src/api/client.test.ts:32:     vi.stubEnv("VITE_API_BASE", "https://api.khelsutra.guru/api");
web/src/api/client.test.ts:37:     vi.stubEnv("VITE_API_BASE",
"https://api.khelsutra.guru/api");
web/src/api/client.ts:20: * set `VITE_API_BASE` at build time to the engine's API base
**including the `/api` path**
```

```
web/src/api/client.ts:25: return (import.meta.env.VITE_API_BASE ?? "/api").replace(/\/+$/, "");
web/src/vite-env.d.ts:9: readonly VITE_API_BASE?: string;
```

11. assistant (2026-06-22T09:33:11.940Z)

`cloudflared.config.example.yml` confirmed (ingress only api.khelsutra.guru ? 127.0.0.1:8000).
`site/scripts/provision-vm.sh` exists. `VITE\_API\_BASE` works as described. Now let me verify the
`provision-vm.sh` systemd unit claim, the `LOCAL\_APPLIANCE\_ARCHITECTURE.md` D1 reference, and
check that `HOSTING\_PLAN.md` is correctly described as the engine's doc (not this repo's). Let
me also check the droplet kit since the doc mentions PR #81/#85 deploy kit alignment.

12. user (2026-06-22T09:33:16.975Z)

```
=== provision-vm systemd unit ===
site/scripts/provision-vm.sh:5:# VM), installs Node, installs + starts a systemd service, opens
the firewall, and VERIFIES
site/scripts/provision-vm.sh:10:# [REMOTE_DIR=/srv/khelsutra-site] [PORT=8787] [NODE_MAJOR=24]
[RESTORE_DB=./backup.db] \
site/scripts/provision-vm.sh:17:DIR="${REMOTE_DIR:-/srv/khelsutra-site}"
site/scripts/provision-vm.sh:40:echo "? [3/5] remote setup (Node $NODE_MAJOR + systemd unit +
firewall) ..."
site/scripts/provision-vm.sh:57:cat > /etc/systemd/system/khelsutra-site.service <<UNIT
site/scripts/provision-vm.sh:66:ExecStart=$NODE_BIN server.js
site/scripts/provision-vm.sh:73:systemctl enable --now khelsutra-site
site/scripts/provision-vm.sh:90:echo "? [5/5] done ? systemd 'khelsutra-site' serving on :$PORT
(systemctl status khelsutra-site)."
=== D1 in LOCAL_APPLIANCE ===
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:30:| D1 | Reverse-proxied **single origin**; engine stays
`127.0.0.1:8000` | `[verified main.py:638]`; 2026-06-20 | CORS `` is moot from the browser ?
the proxy is the boundary |
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:175:**? Alpha deploy path chosen = SPLIT, not the single-
origin shell (owner 2026-06-21).** The alpha ships the **split** topology (SPA on Cloudflare
Pages + engine on an `api.` cloudflared tunnel + CF Access over both ? the engine
`HOSTING_PLAN.md` $2), **not** the single-origin Caddy shell of **T2** (kept as the documented
alternative ? $0/D1). The **deploy kit + runbook**
(`[deploy/DEPLOY_ALPHA.md]`(`../deploy/DEPLOY_ALPHA.md`) +
`deploy/cloudflared.config.example.yml`) is shipped (? templates+runbook); the actual stand-up
(tunnel + Pages + CF Access + key rotation) is the **owner** step. The **T3**
raw-:8000-refused posture test is documented as the deploy kit's $6 verification. Enabled by:
engine exposure-safety boot posture (engine #294) + SPA `VITE_API_BASE` ([#64]).
=== single-origin / Caddy in LOCAL_APPLIANCE ===
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:31:| D2 | **Edge auth**, zero engine auth code |
`HOSTING_PLAN.md`; 2026-06-20 | Caddy basicauth OR cloudflared + Cloudflare Access |
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:54:ONE box: **edge gate** (cloudflared+Access or Caddy
basicauth) ? **reverse proxy / static host** serving the `web/` SPA bundle and proxying `/api/`
? engine **FastAPI `127.0.0.1:8000` ? **LocalStorage** ? **detector seam** (Gemini/stub today;
native WASB on a GPU host). The firewall blocks inbound `8000` + the test ports
`9000/9001/9023`. The marketing site stays a **separate** systemd unit on a **separate**
subdomain.
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:59: ? cloudflared Tunnel + Cloudflare Access (no
inbound ports) OR Caddy :443 + basicauth
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:63:?(2) reverse proxy / static host (Caddy/cloudflared)
?
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:113:**Reused (no re-plumb):** the full `/api` loop; the in-
process worker; the chunked upload endpoints (un-called); LocalStorage default; the detector
seam incl. the landing `NativeWasbRunner`; the **RallyPicker** design + anchors; the `web/`
scaffold; the `HOSTING_PLAN` gate; the `SETUP_NEW_MACHINE` runbook. **Net-new:** Caddyfile +
systemd units; `GET /api/capabilities`; `GET /api/videos`; `GET /api/reels/{video_id}`; `GET
/api/health`; the SPA console screens; ingest-from-bucket wiring; a retention sweep; (reserved)
the bucket-worker bridge.
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:159:| T2 | Appliance shell: Caddyfile + systemd units;
firewall; **edge auth**; rotate key (P1) | ? | owner | ? | ? |
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:175:**? Alpha deploy path chosen = SPLIT, not the single-
origin shell (owner 2026-06-21).** The alpha ships the **split** topology (SPA on Cloudflare
Pages + engine on an `api.` cloudflared tunnel + CF Access over both ? the engine
```

`HOSTING\_PLAN.md` §2), ****not**** the single-origin Caddy shell of ****T2**** (kept as the documented alternative ? §0/D1). The ****deploy kit + runbook**** ([`deploy/DEPLOY\_ALPHA.md`](../deploy/DEPLOY_ALPHA.md) + [`deploy/cloudflared.config.example.yml`](../deploy/cloudflared.config.example.yml)) is shipped (? templates+runbook); the actual stand-up (tunnel + Pages + CF Access + key rotation) is the ****owner**** step. The ****T3**** raw-:8000-refused posture test is documented as the deploy kit's §6 verification. Enabled by: engine exposure-safety boot posture (engine #294) + SPA `VITE\_API\_BASE` ([#64]). docs/LOCAL_APPLIANCE_ARCHITECTURE.md:178:****Blocking gates:**** (1) ****edge-auth**** ? Caddy basicauth (self-contained) vs cloudflared+Cloudflare Access (matches the two-host design, no inbound ports)? (2) ****retention defaults**** ? ****window DECIDED**** (uploads ****7d****, reels durable; [engine #299](https://github.com/Khelsutra/badminton-highlight-indexer/pull/299) tool + scheduling [#301](https://github.com/Khelsutra/badminton-highlight-indexer/issues/301)); ****still open:**** auto-delete originals after compile vs keep proxies+intervals only (privacy/consent). (3) ****per-video Gemini/GPU cost ceiling**** ? what cap, hard-fail vs warn? (4) ****GPU host choice**** ? DO GPU droplet / GCE GPU VM / Cloud Run+GPU (and is a containerized worker worth building for Cloud Run)? (5) ****P5 bridge ownership**** ? who builds the orchestration endpoint + how it coordinates with the active worker sessions.

13. user (2026-06-22T09:33:17.916Z)

=== droplet/config.json.example ===

```
{
  "security": {
    "edge_auth_enabled": true,
    "allowed_video_root": "/srv/khelsutra/incoming",
    "upload_dir": "/srv/khelsutra/incoming/uploads",
    "cf_team_domain": "https://khelsutra.cloudflareaccess.com",
    "cf_access_aud": "PASTE_THE_ACCESS_APPLICATION_AUD_HERE"
  },
  "indexing": {
    "default_segmenter": "gemini"
  }
}
```

=== droplet/engine.env.example ===

KhelSutra engine ? environment for the alpha droplet (the `api.` side).

Copy to /etc/khelsutra/engine.env, fill the secret, and `chmod 600` it.

Loaded by khelsutra-engine.service. Verified against engine origin/master (see ../DEPLOY_ALPHA.md §2 + khelsutra PR #78).

--- Secret (rotate the chat-shared one; Gemini is paid + uploads footage to Google) ---

GEMINI_API_KEY=PASTE_YOUR_ROTATED_KEY_HERE

--- App home: config.json, .env, the SQLite DB and output/ all resolve UNDER here (en

The engine reads \$RALLY_APP_HOME/config.json ? so the security block below liv

/srv/khelsutra/config.json (NOT the repo root). Use `--setup` so writer/reader agree.

RALLY_APP_HOME=/srv/khelsutra

--- Exposure: the two vars the original kit MISSED. Without them the tunnel is dead-on

edge_auth_enabled=true flips the bind to 0.0.0.0 AND keeps the Host-header guard loc

so it boots "healthy" then 400s every request cloudflared forwards with Host: api.khel

RALLY_ALLOWED_HOSTS=api.khelsutra.guru,localhost,127.0.0.1 # REQUIRED ? load-bearing fix

(allowlist is decoupled from the bind); loopback added for on-box health checks

RALLY_BIND_HOST=127.0.0.1

keep the loopback bind (cloudflared

dials 127.0.0.1:8000 anyway)

RALLY_CORS_ALLOW_ORIGINS=https://app.khelsutra.guru # lock CORS to the SPA origin

14. assistant (2026-06-22T09:33:23.532Z)

The droplet kit is consistent with the doc. Now there's a notable discrepancy: the droplet's ``engine.env.example`` sets ``RALLY_ALLOWED_HOSTS=api.khelsutra.guru,localhost,127.0.0.1`` (with loopback added for on-box health checks), while the `DEPLOY_ALPHA.md` §2.2 says ``RALLY_ALLOWED_HOSTS="api.khelsutra.guru"`` (without the loopback entries). This is relevant because §6's posture test does an on-box ``/api/health`` check. Let me note that.

Now let me verify the engine-side claims. The doc makes many claims about the engine repo (``badminton-highlight-indexer``), which is a sibling repo. These are cited as "verified against engine origin/master". Let me spot-check the key falsifiable engine claims since the sibling repo is available locally.

15. user (2026-06-22T09:33:27.222Z)

```
=== engine PR #294 ===
0c0e064 feat(serving): exposure-safety boot posture ? fail fast on a half-configured exposed
engine (default-OFF) (#294)
6fdd1d4 feat(serving): exposure-safety boot posture ? fail fast on a half-configured exposed
engine (default-OFF)
9029415 docs: checkpoint approved structure + sequencing decisions (Config first, atomic
refactor follow-up)
=== test_process_hosted_mode_rejects_nonlocal_uri ===
tests/test_api_endpoints.py:541: Counterweight =
test_process_hosted_mode_rejects_nonlocal_uri (gs:// -> 400 when allowed_video_root
tests/test_api_endpoints.py:608:def test_process_hosted_mode_rejects_nonlocal_uri(env):
=== --setup flag ===
backend/main.py:865: parser.add_argument("--setup", action="store_true", help="Launch
dependency checker and setup utility")
backend/main.py:888: # build system). Resolve its path from this file so `indexer
--setup`
=== --server flag ===
backend/main.py:867: parser.add_argument("--server", action="store_true", help="Run the
FastAPI local API server")
```

16. assistant (2026-06-22T09:33:32.585Z)

The Bash tool ran against the engine repo (cwd reset). Let me confirm which repo that was and verify the remaining engine claims explicitly with absolute paths. The output shows engine PR #294, the ``test_process_hosted_mode_rejects_nonlocal_uri`` test, and ``--setup`/`--server`` flags all exist. Let me verify the ``--setup`` claim more carefully (the doc says ``python -m backend.main --setup`` "targets the exact path the server reads"), the ``[auth]`` extra, and the boot-posture fail-fast behavior.

17. user (2026-06-22T09:33:41.204Z)

```
=== confirm repo ===
origin https://github.com/Khelsutra/badminton-highlight-indexer.git (fetch)
=== [auth] extra in pyproject/setup ===
pyproject.toml:23:authors = [{ name = "Avi Dullu", email = "avi.dullu@gmail.com" }]
pyproject.toml:51: "PyJWT[crypto]>=2.8", # M9 edge-auth: Cf-Access JWT verify (RS256) ?
tests + CI
pyproject.toml:54:# M9 edge-auth (Cloudflare Access). DEFAULT-OFF: only needed when
security.edge_auth_enabled=True.
pyproject.toml:56:auth = ["PyJWT[crypto]>=2.8"]
=== edge_auth_enabled config ===
backend/config/models.py:297: # DEFAULT-OFF: edge_auth_enabled=False ? no JWT is
required/verified (the local single-user
backend/config/models.py:303: edge_auth_enabled: bool = False
backend/config/startup.py:19:behind an edge gate (the owner flips ``security.edge_auth_enabled``
on for the Cloudflare-Access
backend/config/startup.py:30:a strict **no-op** ? exactly the pre-M5 behaviour; with
``edge_auth_enabled=False`` (today's default)
backend/config/startup.py:91: """Exposure-safety posture for an *exposed* HTTP server. Gated
```

```

on ``security.edge_auth_enabled``
backend/config/startup.py:98:     if not sec or not getattr(sec, "edge_auth_enabled", False):
backend/config/startup.py:110:         "security.edge_auth_enabled is ON (engine exposed) but
security.allowed_video_root is "
backend/config/startup.py:124:         f"security.edge_auth_enabled is ON but {missing} is
not configured ? the Cf-Access JWT "
backend/config/startup.py:136:         "security.edge_auth_enabled is ON but PyJWT[crypto] is
not importable ? Cf-Access "
backend/config/startup.py:146:         "security.edge_auth_enabled is ON but
RALLY_ALLOWED_HOSTS is unset ? the Host-header guard "
=== EXPOSED_HOST_ALLOWLIST_LOOPBACK ===
backend/config/startup.py:145:         LEVEL_WARN, "EXPOSED_HOST_ALLOWLIST_LOOPBACK",
=== RALLY_ALLOWED_HOSTS ===
backend/config/startup.py:141:     # loopback-only when RALLY_ALLOWED_HOSTS is unset (see
main._allowed_hosts_from_env) ? it would
backend/config/startup.py:143:     if not (env.get("RALLY_ALLOWED_HOSTS", "") or "").strip():
backend/config/startup.py:146:         "security.edge_auth_enabled is ON but
RALLY_ALLOWED_HOSTS is unset ? the Host-header guard "
backend/config/startup.py:147:         "defaults to loopback-only and will 400 forwarded
requests. Set RALLY_ALLOWED_HOSTS to your "
backend/main.py:89: def _allowed_hosts_from_env(env: Optional[Dict[str, str]] = None) ->
List[str]:
backend/main.py:93:     Host headers by default). ``RALLY_ALLOWED_HOSTS`` (comma-separated)
overrides; the default keeps
backend/main.py:98:     raw = (src.get("RALLY_ALLOWED_HOSTS", "") or "").strip()
backend/main.py:111:     allowed_hosts=_allowed_hosts_from_env(),
backend/main.py:1088:         # the Host allowlist (the default; don't set
RALLY_ALLOWED_HOSTS to a non-loopback-only set
backend/main.py:1101:         f"RALLY_ALLOWED_HOSTS to your public host(s) (the
Host-header guard defaults to "

```

18. assistant (2026-06-22T09:33:48.691Z)

Excellent. The engine claims hold up strongly:

- `[auth]` extra = `PyJWT[crypto]>=2.8` confirmed (pyproject.toml:56)
- `edge\_auth\_enabled` default False, fail-fast posture confirmed (startup.py)
- `EXPOSED\_HOST\_ALLOWLIST\_LOOPBACK` warning confirmed (startup.py:145)
- The `RALLY\_ALLOWED\_HOSTS` Host-header guard behavior confirmed

Let me verify the remaining key claims: the bind-flip to 0.0.0.0 when edge_auth on, the fail-fast (raises) vs warn distinction, the 401 (not 500) for missing PyJWT, and the boot posture failing on missing jail/cf_team_domain/cf_access_aud.

19. user (2026-06-22T09:33:50.514Z)

```

=== bind flip to 0.0.0.0 when edge_auth on ===
backend/main.py:116: def _resolve_bind_host(config: Optional[AppConfig] = None, cli_host:
Optional[str] = None) -> str:
backend/main.py:117:     """P0c: the server bind address. Precedence: ``--host`` >
``RALLY_BIND_HOST`` env > ``0.0.0.0``
backend/main.py:123:     env = os.environ.get("RALLY_BIND_HOST")
backend/main.py:128:     return "0.0.0.0" # behind a Cloudflare-Access-style edge gate (M9)
? bind all interfaces
backend/main.py:873:         help="Bind address. Default 127.0.0.1 (localhost-
only); 0.0.0.0 when "
backend/main.py:1082:         bind_host = "127.0.0.1"
backend/main.py:1083:         port = _resolve_app_port(bind_host, args.port)
backend/main.py:1086:         url = f"http://{bind_host}:{port}/"
backend/main.py:1095:         bind_host = _resolve_bind_host(global_config, args.host) #
P0c: config/profile-driven (default 127.0.0.1)
backend/main.py:1097:         print(f"[SERVER] Starting local FastAPI server on
http://{bind_host}:{port}...")
=== fail-fast: raises on missing jail/aud ===

```

```

out: list[StartupCheck] = []

```

```

# (1) The path-jail is load-bearing once authenticated tenants can hit /api/process: WITHOUT
it,
# a tenant can point the engine at an ARBITRARY local file (allowed_video_root unset = "any
local
# file", the single-user default). Fatal for an exposed engine.
root = getattr(sec, "allowed_video_root", None)
if not root or not str(root).strip():
    out.append(StartupCheck(
        LEVEL_ERROR, "EXPOSED_NO_VIDEO_ROOT",
        "security.edge_auth_enabled is ON (engine exposed) but security.allowed_video_root
is "
        "unset ? an authenticated tenant could make /api/process read an arbitrary local
file. "
        "Set allowed_video_root to a jail directory before exposing the engine.",
    ))

# (2) Edge auth cannot verify a token without the team domain + AUD. Today that only fails
on the
# FIRST request (auth.resolve_tenant); lift it to BOOT so a misconfigured gate never starts
# "healthy" then 401s every request.
team = (getattr(sec, "cf_team_domain", "") or "").strip()
aud = (getattr(sec, "cf_access_aud", "") or "").strip()
if not team or not aud:
    missing = " + ".join(n for n, v in (("cf_team_domain", team), ("cf_access_aud", aud)) if
not v)
    out.append(StartupCheck(
        LEVEL_ERROR, "EXPOSED_EDGE_AUTH_INCOMPLETE",
        f"security.edge_auth_enabled is ON but {missing} is not configured ? the Cf-Access
JWT "
        "cannot be verified (every request would fail). Set security.cf_team_domain "
        "(e.g. https://<team>.cloudflareaccess.com) and security.cf_access_aud (the CF
Access "
        "application AUD tag).",
    ))

# (3) The auth extra (PyJWT[crypto]) is lazy-imported by auth.verify_cf_access_token. If
it's
# missing, every request 500s at verify time ? a boot WARN beats a per-request 500. Warning,
not
# fatal: the auth module's own ImportError is the final authority.
if not _auth_extra_available():
    out.append(StartupCheck(
        LEVEL_WARN, "EXPOSED_AUTH_EXTRA_MISSING",
        "security.edge_auth_enabled is ON but PyJWT[crypto] is not importable ? Cf-Access "
        "verification will fail at request time. Install the 'auth' extra (pip install -e
.[auth]).",
    ))

# (4) P0c: an exposed engine binds 0.0.0.0, but the Host-header (DNS-rebinding) guard
defaults to
# loopback-only when RALLY_ALLOWED_HOSTS is unset (see main._allowed_hosts_from_env) ? it
would
# then 400 every request the edge forwards with its public Host. Boot WARN beats a silent
400-all.
if not (env.get("RALLY_ALLOWED_HOSTS", "") or "").strip():
    out.append(StartupCheck(
        LEVEL_WARN, "EXPOSED_HOST_ALLOWLIST_LOOPBACK",
        "security.edge_auth_enabled is ON but RALLY_ALLOWED_HOSTS is unset ? the Host-header
guard "
        "defaults to loopback-only and will 400 forwarded requests. Set RALLY_ALLOWED_HOSTS
to your "
        "public host(s) (or '*' when the edge gate validates the host).",
    ))

return out

```

```
def run_startup_checks(
    config: Any,
    *,
    cuda_available: Optional[bool] = None,
    env: Optional[Mapping[str, str]] = None,
    include_exposure: bool = False,
) -> list[StartupCheck]:
```

20. assistant (2026-06-22T09:33:56.177Z)

All engine claims verified precisely:

- Bind flip to `0.0.0.0` when edge_auth on, overridable by `RALLY\_BIND\_HOST` (main.py:117,128) ? confirmed
- Fail-fast `LEVEL\_ERROR` on missing `allowed\_video\_root` (EXPOSED_NO_VIDEO_ROOT) and missing `cf\_team\_domain`/`cf\_access\_aud` (EXPOSED_EDGE_AUTH_INCOMPLETE) ? confirmed
- `WARN` (not fatal) for missing PyJWT ? confirmed

One nuance: the doc §2.1 says missing PyJWT means "every authenticated request is rejected 401". The engine comment (startup.py line ~135) says "every request 500s at verify time". The doc's changelog claims it fixed "401-not-500 wording (§2.1)" ? but the engine code's own comment says 500. Let me check the actual auth module behavior to see whether it's 401 or 500 when PyJWT is missing.

21. user (2026-06-22T09:33:59.352Z)

=== auth module: missing PyJWT behavior ===

backend/api/auth.py

backend/config/startup.py

backend/main.py

--- find auth file ---

backend/api/auth.py

=== ImportError / 401 / 500 in auth ===

backend/api/auth.py:31: """A required Cf-Access token was missing or failed verification ? the caller returns 401/403."""

backend/api/auth.py:107: Enabled ? require + verify the Cf-Access header; raise `EdgeAuthError` (? 401/403) on any miss.

22. user (2026-06-22T09:34:04.533Z)

=== auth.py PyJWT import + error handling ===

"""Cloudflare Access edge-auth ? verify the Cf-Access JWT, derive the tenant `user\_id` (M9, D9).

DEFAULT-OFF: when `security.edge\_auth\_enabled` is False the resolver is a strict NO-OP (returns `None` = the local single-user) ? nothing about the existing local path changes. When enabled, the `Cf-Access-Jwt-Assertion` header is **required** and **verified** (RS256) against the team JWKS ? signature + exp + issuer + audience ? so a direct hit to the Cloud Run URL **can't spoof** the identity header (the explicit D9 anti-spoof requirement). Returns the verified `user\_id` (the token's `email`, else `sub`).

`jwt` (`PyJWT[crypto]`) is **lazy-imported** so the base runtime works without it (install the `auth` extra for hosted/multi-tenant). The JWKS fetch is **injectable**, so the whole verifier is unit-tested **offline** with an in-test RSA keypair + a fake JWKS ? no network, no Cloudflare.

from __future__ import annotations

import json

import time

```
import urllib.request
from typing import Any, Callable, Mapping, Optional
```

Cloudflare Access sends the signed assertion in this header (it terminates at the CF edge)

```
CF_ACCESS_HEADER = "Cf-Access-Jwt-Assertion"
_JWKS_PATH = "/cdn-cgi/access/certs"
_JWKS_TTL_SEC = 600.0 # re-fetch the team JWKS at most every 10 min (keys rotate slowly)
```

```
JwksFetcher = Callable[[str], dict]
```

```
class EdgeAuthError(Exception):
    """A required Cf-Access token was missing or failed verification ? the caller returns
    401/403."""
```

module-level JWKS cache: team_domain -> (fetched_at, jwks_dict). Production only; tests

```
_jwks_cache: dict[str, tuple[float, dict]] = {}
```

```
def _issuer(team_domain: str) -> str:
    return team_domain.rstrip("/")
```

```
def _default_jwks_fetcher(team_domain: str) -> dict:
    """Fetch the team's public JWKS (`<team_domain>/cdn-cgi/access/certs`). Network ? only the
    production path reaches here; tests pass an injected fetcher."""
    url = _issuer(team_domain) + _JWKS_PATH
    with urllib.request.urlopen(url, timeout=5) as resp: # noqa: S310 (https team domain)
        data: Any = json.load(resp)
    if not isinstance(data, dict) or "keys" not in data:
        raise EdgeAuthError(f"unexpected JWKS shape from {url}")
    return data
```

```
def _get_jwks(team_domain: str, fetcher: Optional[JwksFetcher]) -> dict:
    """Return the JWKS for `team_domain`, cached with a TTL. An injected `fetcher` (tests)
    bypasses the cache so each test is deterministic."""
    if fetcher is not None:
        return fetcher(team_domain)
    now = time.time()
    cached = _jwks_cache.get(team_domain)
    if cached and (now - cached[0]) < _JWKS_TTL_SEC:
        === how EdgeAuthError maps to status ===
backend/api/auth.py:30:class EdgeAuthError(Exception):
backend/api/auth.py:49:         raise EdgeAuthError(f"unexpected JWKS shape from {url}")
backend/api/auth.py:69:         return the user_id (`email` else `sub`). Raises `EdgeAuthError`
on ANY failure."""
backend/api/auth.py:74:         raise EdgeAuthError(
backend/api/auth.py:82:             raise EdgeAuthError("token header has no kid")
backend/api/auth.py:89:     except EdgeAuthError:
backend/api/auth.py:92:         raise EdgeAuthError(f"no JWKS key for kid {exc}") from exc
backend/api/auth.py:94:         raise EdgeAuthError(f"Cf-Access token verification failed:
{exc}") from exc
backend/api/auth.py:98:         raise EdgeAuthError("verified token carries no email/sub claim")
backend/api/auth.py:107:     Enabled ? require + verify the Cf-Access header; raise
`EdgeAuthError` (? 401/403) on any miss.
```

23. user (2026-06-22T09:34:08.421Z)

```
=== PyJWT lazy import + what raises ===
    return the user_id (`email` else `sub`). Raises `EdgeAuthError` on ANY failure."""
    try:
```

```

import jwt
from jwt import PyJWKSet
except Exception as exc: # pragma: no cover - only when the auth extra isn't installed
    raise EdgeAuthError(
        "PyJWT[crypto] is required for edge auth ? install the 'auth' extra "
        "(pip install -e .[auth]) on a host with edge_auth_enabled."
    ) from exc

try:
    kid = jwt.get_unverified_header(token).get("kid")
    if not kid:
        raise EdgeAuthError("token header has no kid")
    signing_key = PyJWKSet.from_dict(jwks)[kid]
    claims = jwt.decode(
        token, signing_key.key, algorithms=["RS256"],
        audience=aud, issuer=_issuer(team_domain),
        options={"require": ["exp", "iss", "aud"]},
    )
except EdgeAuthError:
    raise
except KeyError as exc: # kid not in the JWKS
    raise EdgeAuthError(f"no JWKS key for kid {exc}") from exc
except Exception as exc: # jwt.* errors: bad sig / expired / wrong aud|iss / malformed
    raise EdgeAuthError(f"Cf-Access token verification failed: {exc}") from exc

user = claims.get("email") or claims.get("sub")
if not user:
    raise EdgeAuthError("verified token carries no email/sub claim")
return str(user)

```

24. assistant (2026-06-22T09:34:18.854Z)

Confirmed: when PyJWT is missing, the code raises `EdgeAuthError`, which the docstring (line 31, 107) maps to "401/403". So the doc's §2.1 claim ("else every authenticated request is rejected 401") is CORRECT against the current engine code. The engine's own startup.py comment still says "500s" but that's a stale comment in the engine, not relevant to this doc ? the actual auth-module behavior backs the doc's 401 wording.

Now let me address the one real discrepancy I found: the §6 on-box health check. The DEPLOY_ALPHA.md §2.2 sets `RALLY\_ALLOWED\_HOSTS="api.khelsutra.guru"` (only the public host), but §6's posture test and the §2.3/§6 health checks include `curl -i https://api.khelsutra.guru/api/health`. The droplet kit's `engine.env.example` adds `localhost,127.0.0.1` precisely so on-box health checks (loopback Host) don't 400. Let me verify whether the §6 verify steps actually use a loopback Host that would be blocked.

Looking at §6: the external check is `http://<box-public-ip>:8000/api/health` (expected refused), and `https://api.khelsutra.guru/api/health` (via tunnel, Host = api.khelsutra.guru ? covered by the doc's allowlist). So §6 doesn't strictly need loopback. The droplet kit adds loopback as a hardening extra. This is a minor inconsistency between the runbook and its own kit, not a break. I'll note it as low/medium.

Let me do a final check on a couple remaining specifics: the `max\_part\_mb` default of 95, and the `cf\_team\_domain`/`cf\_access\_aud` config field names.

25. user (2026-06-22T09:34:21.084Z)

```

=== max_part_mb default ===
backend/config/models.py:293:    max_part_mb: int = 95
=== cf_team_domain / cf_access_aud / upload_dir / allowed_video_root config fields ===
backend/config/models.py:128:    default_segmenter: str = "yolo_hybrid"
backend/config/models.py:210:    @field_validator("default_segmenter")
backend/config/models.py:283:    allowed_video_root: Optional[str] = None
backend/config/models.py:287:    # arrive with PR-3 identity scoping. ? In hosted mode, set
allowed_video_root to
backend/config/models.py:288:    # contain upload_dir or /api/process will reject the uploaded

```

```

paths.
backend/config/models.py:289:    upload_dir: str = "output/uploads"
backend/config/models.py:304:    cf_team_domain: str = ""          # e.g.
https://<team>.cloudflareaccess.com (issuer + /certs JWKS)
backend/config/models.py:305:    cf_access_aud: str = ""          # the CF Access application AUD
tag the token must carry

```

26. assistant (2026-06-22T09:34:29.719Z)

All verified:

- `max\_part\_mb` default = 95 (models.py:293) ? matches doc §7 "default 95"
- `default\_segmenter` ships as `yolo\_hybrid` (models.py:128) ? matches doc §2.2 "the shipped default `yolo\_hybrid` needs torch/GPU"
- `cf\_team\_domain`, `cf\_access\_aud`, `upload\_dir`, `allowed\_video\_root` config field names all confirmed

The doc is remarkably well-verified and was just refreshed in the 2026-06-22 changelog. The status `READY` is accurate (owner-executed cloud steps remain, which is correctly framed as owner-gated). PRs #64, #65, #78 all confirmed. This is NOT an archive candidate ? it's a live runbook whose owner-execution steps haven't all completed.

One thing to double-check: MEMORY says `app.khelsutra.guru` is "live + Access-gated" as of 2026-06-22 (T2 alpha deploy DONE). The doc's status is still `READY` ("owner executes the cloud steps") rather than reflecting that the deploy is now LIVE. This is a mild staleness ? the doc describes the stand-up as pending/owner-to-do, but per the session memory the alpha is actually deployed and live. However, that "live" state is from session memory, not verifiable in this repo. Let me check if the repo itself records the deploy as live anywhere.

27. user (2026-06-22T09:34:33.811Z)

```

=== repo mentions of app.khelsutra.guru live/deployed ===
deploy/DEPLOY_ALPHA.md
deploy/droplet/README.md
=== droplet README live status ===
deploy/droplet/README.md:100:> **? Live-verified on the droplet (2026-06-22).** Engine installed
+ booted on `168.144.126.176`; full
deploy/droplet/README.md:103:> exercised live. Tunnel + Access + Pages are up; the full **DNS ?
Cloudflare ? Access ? tunnel ? engine**
deploy/droplet/README.md:108:> **? Two gotchas the live run surfaced:**
=== LOCAL_APPLIANCE T-row status for alpha deploy ===
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:23:**Outcome ("good"):** an operator signs in, picks a
source, watches it index, curates rallies, downloads a reel ? on the **live CPU droplet today**
(Gemini/stub), and on a **GPU host later** (native WASB) with **no UI fork**.
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:34:| D5 | Console lives in the khelsutra **`web/`** React
scaffold | `web/README.md`; 2026-06-20 | Cloudflare-Pages-style root, not engine static |
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:39:| D10 | Delivery = a **self-hosted webserver + UI the
user spins up** (NOT a desktop app); the user points at **local storage + GPU/CPU OR cloud
locations** and the UI drives the tool | Owner 2026-06-20 | **supersedes
`DESKTOP_APP_PLAN.md`**; one artifact (this appliance) runs on the owner box, a droplet, or a
GPU host ? same UI, compute/storage are settings |
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:58:    ? (1) EDGE GATE ? the ONLY ingress; auth lives HERE,
not in the app
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:88:| **Google Drive** (`gdrive://`) | `parse_uri`
recognizes it, but the gdrive backend is an **emulated stub** `[verified ref.py,
DECOUPLED_COMPUTE research]` | **stub ? roadmap;** MVP requires a Drive?bucket pre-step; do
**not** present Drive as live until the backend is real |
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:103:No engine auth, by design; the gate is external.
**Invariant (all three must hold):** engine bound to `127.0.0.1` + firewall closes `8000` + edge
gate up. Go-live checklist `[HOSTING_PLAN.md:102-108]`: Access/basicauth ON; **rotate the chat-
shared `GEMINI_API_KEY`**; set `security.allowed_video_root`; rate-limit; per-user quota.
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:123:- **Log-analysis pass before merge (DoD item).** Every
PR that adds a code path includes a check that the path logs meaningfully, and the **logs are
inspected during CUJ live-replay** (§8 Layers 1?3). Ties [[feedback-launch-quality-bar]].
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:128:| No app auth + CORS `` `[main.py:42]` | open if any
of {localhost-bind, firewall, gate} slips | enforce all three; **LIVE posture test** that raw

```

```
`:8000` is refused from outside |  
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:134:| Drive ingest overstated | `gdrive://` is a stub |  
don't present Drive as live; require Drive?bucket pre-step until real |  
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:142:## 8. CUJs & live-test plan `[required]`
```

28. assistant (2026-06-22T09:34:44.595Z)

Important finding. The sibling doc `deploy/droplet/README.md:100` records "? Live-verified on the droplet (2026-06-22). ... Tunnel + Access + Pages are up; the full DNS ? Cloudflare ? Access ? tunnel ? engine ... chain ... exercised live." So the alpha stand-up has actually been DONE and is LIVE.

But `DEPLOY\_ALPHA.md`'s status header still reads `READY` ("owner executes the cloud steps") and \$3/\$4/\$5 are written entirely in the imperative future tense as if the cloud stand-up has not happened yet. This is a genuine staleness: a sibling doc in the same repo, dated the same day (2026-06-22), says the deploy is live-verified, while this runbook still presents it as not-yet-executed. A reader would conclude the stand-up is still pending. This is medium severity (the runbook is still useful as a procedure, but its top-line status is now behind reality).

Let me read that droplet README block in full to quote it precisely.

29. user (2026-06-22T09:34:45.273Z)

```
96      **The shareable CUJ:** open `app.khelsutra.guru` ? sign in (email OTP) ? **upload a  
video** (lands in the  
97      jail) ? Process ? pick rallies ? Build ? download the reel. *(Pasting a `gs://` URI is  
rejected by the  
98      jail on the exposed box ? upload is the alpha path; see `../DEPLOY_ALPHA.md` $7.)*  
99  
100     > **? Live-verified on the droplet (2026-06-22).** Engine installed + booted on  
`168.144.126.176`; full  
101     > suite **1313 passed** + ruff/mypy clean on the box; boot-posture fail-fast, the Host-  
guard `400?200`,  
102     > the path-jail (`gs://`/etc/passwd`/`...` ? 400), the T4 upload?jail, and the edge-  
auth `401`-gate all  
103     > exercised live. Tunnel + Access + Pages are up; the full **DNS ? Cloudflare ? Access ?  
tunnel ? engine**  
104     > chain is confirmed (`/api/health` = 200 through an Access session; unauthenticated ?  
302 to the login).  
105     > Two things still need a human: the rotated **Gemini key** (`/etc/khelsutra/engine.env`  
? restart) and a  
106     > **browser one-time-PIN login** to exercise the upload?reel CUJ.  
107     >  
108     > **? Two gotchas the live run surfaced:**  
109     > - **`indexing.default_segmenter` MUST be `gemini` on a CPU box** (set in  
`config.json.example`). The SPA  
110     > sends no segmenter, so the engine uses its default ? the shipped default  
`yolo_hybrid` needs torch/GPU  
111     > and every job fails `No module named 'torch'` without this.  
112     > - **Without a valid `GEMINI_API_KEY`, a job finishes `done` with 0 rallies** ? the  
missing-key error lands  
113     > only in the engine log, not the job status. Set the key before sharing.  
114
```

30. assistant (2026-06-22T09:34:55.895Z)

This confirms the staleness conclusively. The droplet README (same repo, same date) says the alpha is "Live-verified on the droplet (2026-06-22)" with the full chain confirmed and \$6's exact posture tests already exercised live. Yet `DEPLOY\_ALPHA.md` still reads `Status: READY` (owner executes the cloud steps) and \$6 is written as a "Verify (...)" to-do rather than a passed result.

I have enough to render the verdict. Let me compile the findings.

Summary of verification:

- ****All file/path/route references****: confirmed (``deploy/cloudflared.config.example.yml``, ``docs/LOCAL_APPLIANCE_ARCHITECTURE.md`` D1, ``site/scripts/provision-vm.sh`` systemd unit, ``web/VITE_API_BASE`` in `client.ts`).
- ****All PR numbers****: #64 (VITE_API_BASE), #65 (deploy kit), #294 (engine boot posture) all confirmed.
- ****All engine config/behavior claims****: confirmed against engine origin/master (edge_auth fail-fast, EXPOSED_HOST_ALLOWLIST_LOOPBACK warn, bind-flip to 0.0.0.0, max_part_mb=95, yolo_hybrid default, [auth]=PyJWT, gs:// jail test, 401-on-missing-PyJWT).
- ****Staleness****: the top-line ``READY`` status and the future-tense §3-§6 are now behind reality ? the same-repo droplet README records the deploy as live-verified on 2026-06-22.
- ****Minor inconsistency****: §2.2 sets ``RALLY_ALLOWED_HOSTS="api.khelsutra.guru"`` only, while the droplet ``engine.env.example`` adds ``localhost,127.0.0.1`` for on-box health checks.

31. user (2026-06-22T09:35:34.317Z)

Output does not match required schema: root: must have required property 'findings'

32. assistant (2026-06-22T09:35:37.331Z)

The ``findings`` array must be a top-level property; it seems the nested object structure broke. Let me re-call with the correct shape.

33. user (2026-06-22T09:36:08.905Z)

Output does not match required schema: root: must have required property 'findings'

34. assistant (2026-06-22T09:36:15.484Z)

The schema validator reports ``findings`` missing, which suggests the large JSON string in ``findings`` may be getting parsed/truncated. Let me try with a more compact findings array (fewer, shorter entries) to ensure it parses cleanly.

35. user (2026-06-22T09:36:40.606Z)

Output does not match required schema: root: must have required property 'findings'

36. assistant (2026-06-22T09:36:44.833Z)

The validator keeps reporting ``findings`` missing despite it being present. The issue may be that the JSON array string is being passed but the harness expects an actual array type, or one of my entries has a parsing issue. Let me try with a drastically reduced findings array (just 3 entries) to isolate whether it's a size/parsing issue.

37. user (2026-06-22T09:36:50.640Z)

Output does not match required schema: root: must have required property 'findings'

38. assistant (2026-06-22T09:36:56.380Z)

Even a tiny array fails with the same error. The validator says ``findings`` is missing at root, which means the harness is likely interpreting my JSON-string value as a string, not an array ? or expecting ``findings`` as a native array argument rather than a stringified one. The schema declares ``findings`` as ``type: array``. My value is a JSON string. Let me pass it without the outer quotes issue ? the tool input must receive an actual array. The most likely cause: my findings value is being sent as a string. Let me retry, ensuring findings is structured as an array of objects natively (the framework should accept the array literal).

39. user (2026-06-22T09:36:58.820Z)

Output does not match required schema: root: must have required property 'findings'

40. assistant (2026-06-22T09:37:04.163Z)

Still failing even with a single well-formed entry. The error is at `root` level claiming `findings` is required and absent ? which means the harness is not seeing my `findings` parameter at all, possibly because it's being passed as a string and the schema wants an array but the binding drops it. Let me check what the actual valid keys are by reading the error literally: "root: must have required property 'findings'".

The `doc`, `overallStatus`, `summary` are clearly accepted (no complaint about them). Only `findings` is reported missing. This strongly suggests the `findings` value I'm sending is being rejected at the type level (string vs array) and thus treated as absent.

Let me try passing findings as an empty array literal first to confirm the binding works, then build up.

41. user (2026-06-22T09:37:05.208Z)

Structured output provided successfully

42. assistant (2026-06-22T09:37:12.414Z)

An empty `[]` works, but a populated array does not. That means one of my finding objects violates the per-item schema, and the validator surfaces it as the top-level `findings` error. The item schema requires exactly: `claim`, `docLocation`, `verdict`, `evidence`, `suggestedFix`, `severity` ? and `verdict` enum is `["stale","correct","unverifiable"]`, `severity` enum is `["high","medium","low"]`. My single-item attempt had all those.

Wait ? the schema item lists `required: [claim, docLocation, verdict, evidence, suggestedFix, severity]` but NOT all properties may be allowed if `additionalProperties:false`. My items only use allowed keys. The single valid-looking item still failed though.

The difference between the `[]` success and the failures is purely the presence of content. Possibly the harness needs the array passed as a real JSON value and my multi-line/escaped strings are breaking the outer parse. Let me try one item with minimal, simple ASCII strings and no special characters.

43. user (2026-06-22T09:37:13.470Z)

Structured output provided successfully

44. assistant (2026-06-22T09:37:18.072Z)

A simple populated item works. The earlier failures were caused by a special character in my content breaking the JSON string parse ? most likely the `>` arrows, `<`, or the `:` / quotes inside evidence strings. Let me rebuild the full findings using only safe ASCII (no `>`, replace with "to"; no `<`; avoid embedded double quotes). I'll keep all 14 findings but sanitize the strings.

45. user (2026-06-22T09:37:47.530Z)

Output does not match required schema: root: must have required property 'findings'

46. assistant (2026-06-22T09:37:54.376Z)

The full array fails again. Since a single simple item works, the breakage is content in one of the larger items. The likely culprit: `>` in the auth finding (`PyJWT[crypto]>=2.8`) or `[auth]` / `[crypto]` square brackets confusing the parser, or `feat(web)` parens. Let me bisect: submit the first 7 findings only, sanitizing `>` to "ge" and removing inner square brackets.

47. user (2026-06-22T09:38:11.597Z)

Output does not match required schema: root: must have required property 'findings'

48. assistant (2026-06-22T09:38:14.489Z)

Even the first 7 (sanitized) fail. Yet a single item works. The breakage must be a size/length threshold or a specific item among 2-7. Let me bisect to 3 items.

49. user (2026-06-22T09:38:19.989Z)

Output does not match required schema: root: must have required property 'findings'

50. assistant (2026-06-22T09:38:26.079Z)

That confirms it: when I OMIT `findings` entirely, I get the identical "must have required property findings" error. So the earlier failures were NOT my content ? they were that the `findings` parameter wasn't being transmitted/parsed as present in those calls. The empty `[]` and single-item calls that "succeeded" prove the binding works with small payloads, but the large multi-item payloads are being dropped (likely the value exceeds a length limit and gets truncated, so the parameter is lost).

So this is a payload-size problem. I need to keep the total `findings` JSON small enough to transmit. Let me drastically shorten each finding (terse evidence, short claims) and include all of them compactly.

51. user (2026-06-22T09:38:46.766Z)

Output does not match required schema: root: must have required property 'findings'

52. assistant (2026-06-22T09:38:49.858Z)

Still failing at 14 compact items. The size threshold is lower than this. Let me find the working size by submitting 6 compact items.

53. user (2026-06-22T09:39:03.306Z)

Output does not match required schema: root: must have required property 'findings'

54. assistant (2026-06-22T09:39:05.821Z)

7 items still fails. The single-item call worked. Let me bisect precisely: try 4 items with short content.

55. user (2026-06-22T09:39:12.876Z)

Output does not match required schema: root: must have required property 'findings'

56. assistant (2026-06-22T09:39:15.381Z)

4 items fails; 1 item works. Let me try 2 items.

57. user (2026-06-22T09:39:18.534Z)

Output does not match required schema: root: must have required property 'findings'

58. assistant (2026-06-22T09:39:21.064Z)

2 items fails, 1 works. So the limit is effectively 1 item per call, OR there's something specific. Let me test 2 trivial items to confirm it's a count/size issue rather than content.

59. user (2026-06-22T09:39:23.117Z)

Structured output provided successfully

60. assistant (2026-06-22T09:39:30.920Z)

Two trivial items work. So it IS a payload-size/length issue, not item count ? the longer string values are pushing the serialized argument past a transmission/truncation limit, causing

`findings` to be dropped. I need to keep total content small. Let me build the full set of findings but with very terse strings, staying under the size that broke (the 2-item call with longer strings failed). I'll consolidate the 12 "correct" findings into a couple of grouped entries and keep evidence ultra-short.

61. user (2026-06-22T09:39:42.781Z)

Output does not match required schema: root: must have required property 'findings'

62. assistant (2026-06-22T09:39:47.288Z)

5 grouped items still too large. The threshold is very tight. Let me cut to exactly the 2 staleness findings plus 1 grouped "correct" plus the unverifiable ? 4 items, but with shorter strings than my failing 4-item attempt. Actually my earlier 4-item attempt failed too. The working ceiling appears to be quite low. Let me reduce to 3 items with terse strings.

63. user (2026-06-22T09:39:53.835Z)

Output does not match required schema: root: must have required property 'findings'

64. assistant (2026-06-22T09:39:57.671Z)

3 items failed, but this time the `summary` is much longer (I moved the correct-claims list into summary). The failure may be driven by the TOTAL argument size including `summary`. Let me shorten the summary too and keep 3 terse findings.

65. user (2026-06-22T09:40:02.140Z)

Output does not match required schema: root: must have required property 'findings'

66. assistant (2026-06-22T09:40:09.384Z)

3 terse items still fail; 2 trivial items worked. The boundary is between 2 and 3 items, or a per-call total just above the 2-trivial-item size. This is a hard transmission ceiling. I'll deliver the findings across what fits ? but the task requires ONE call. Given a single call must carry the answer and the harness truncates above a small size, I must fit everything into the smallest possible payload. Let me try 3 ultra-minimal items (matching the trivial-item character budget).

67. user (2026-06-22T09:40:13.243Z)

Output does not match required schema: root: must have required property 'findings'

68. assistant (2026-06-22T09:40:17.417Z)

Even 3 ultra-minimal items fail, while 2 trivial worked. The ceiling is exactly 2 items in practice. This is clearly a harness transmission cap, not my content. I must convey all material findings within 2 array entries. I'll merge: entry 1 = the two staleness issues consolidated, entry 2 = the broad "everything else verified correct + cost unverifiable" summary. I'll keep both terse.

69. user (2026-06-22T09:40:28.304Z)

Output does not match required schema: root: must have required property 'findings'

70. assistant (2026-06-22T09:40:33.227Z)

2 items but with long strings ? failed. The 2 that worked were trivial (single-char values). So it really is a total-character budget, not item count. I need both items to be short. Let me trim aggressively to fit roughly the trivial-2-item budget while preserving the essential signal.

71. user (2026-06-22T09:40:38.281Z)

Output does not match required schema: root: must have required property 'findings'

72. assistant (2026-06-22T09:40:41.435Z)

Still failing. The working budget is genuinely tiny. Let me match almost exactly the size of the successful 2-trivial-item call by keeping values to a few words each.

73. user (2026-06-22T09:40:44.480Z)

Output does not match required schema: root: must have required property 'findings'

74. assistant (2026-06-22T09:40:49.461Z)

This is surprising ? even this short 2-item call failed, while the earlier 2-trivial-item one (single letters) succeeded. The variable isn't pure length then. The difference: my failing items contain characters like `:`, `.` , `/` , `,` , `;` , `#` , `-` inside string values. The successful trivial items used only single letters. A specific character may be breaking the argument parse. The `#` (in `#294`) and `/` and `:` are prime suspects. Let me test 2 items that contain colons and slashes but no other punctuation, to isolate.

75. user (2026-06-22T09:40:52.016Z)

Structured output provided successfully

76. assistant (2026-06-22T09:40:58.172Z)

That worked ? colons and slashes are fine. So the breaking element is a specific character I'd been using: most likely the comma inside values (e.g. `localhost,127.0.0.1`), the semicolon, or the `#`. Actually this successful call had a colon/slash and a `:3` line ref and worked. The difference from the immediately prior failure: the prior had `;` and `#` and `,`. Let me now build the real 2-item answer using only safe punctuation (colons, slashes, periods, spaces, hyphens) ? avoiding commas, semicolons, and `#` inside string values.

77. user (2026-06-22T09:41:11.907Z)

Output does not match required schema: root: must have required property 'findings'

78. assistant (2026-06-22T09:41:16.338Z)

Back to failure with the longer content. So it IS a size cap after all ? the probe that worked was short. The cap is around the size of two short items. I'll deliver the final answer as two reasonably-short items that still carry the essential, actionable findings, trimming evidence to the most load-bearing pointers.

79. user (2026-06-22T09:41:23.990Z)

Output does not match required schema: root: must have required property 'findings'

80. assistant (2026-06-22T09:41:27.106Z)

Still failing. The boundary is frustratingly tight. The last confirmed SUCCESS had item strings of ~4-6 words each. Let me match that exact scale precisely ? short phrases, no long evidence chains.

81. user (2026-06-22T09:41:30.577Z)

Output does not match required schema: root: must have required property 'findings'

82. assistant (2026-06-22T09:41:36.070Z)

Even this short version fails, yet my probe with "alpha beta gamma delta" / "droplet README line 100 live" succeeded. The ONLY material difference now is punctuation inside the `evidence`/`docLocation`: `deploy/DEPLOY\_ALPHA.md:3` has a colon, my probe used

`deploy/DEPLOY\_ALPHA.md:3` too and worked. The difference: `2026-06-22` (hyphens + digits) and `:100`. My probe didn't have those. Let me test whether a date string like `2026-06-22` or `README.md:100` breaks it.

83. user (2026-06-22T09:41:38.550Z)

Structured output provided successfully

84. assistant (2026-06-22T09:41:47.269Z)

That succeeded ? so dates, colons, line refs all work. The breaking factor in the immediately-prior failure was the longer `summary` ("All claims verified correct against the live repo...") combined with the items. It really is a cumulative size cap, and I'm right at the edge. The successful call had a SHORT summary ("probe date and line refs").

So the constraint: keep `summary` short AND keep findings to 2 compact items. The above call actually captured both staleness findings successfully. I now want to add the "everything else verified correct" coverage. Given the cap allows ~2 short items + short summary, let me make my final answer fit: short summary, and 2 findings where I fold the broad "correct" coverage into the summary text (which is allowed to be moderately sized) rather than a third finding. Let me find how much summary I can add by testing with the two real staleness findings + a richer summary.

85. user (2026-06-22T09:41:53.216Z)

Output does not match required schema: root: must have required property 'findings'

86. assistant (2026-06-22T09:42:01.249Z)

That richer summary tipped it over. The cap is on TOTAL serialized size. The prior success had a tiny summary. I'll keep the summary moderate and the two findings compact ? capturing the verdict (minor-staleness), the two staleness items with fixes, and a concise statement that the rest verified correct. Let me trim the summary to a middle size and re-test.

87. user (2026-06-22T09:42:04.751Z)

Output does not match required schema: root: must have required property 'findings'

88. assistant (2026-06-22T09:42:10.446Z)

That summary is still too long. The successful one had summary "probe date and line refs" (~5 words). The cap is extremely tight on total payload. I'll keep summary very short and rely on the two findings to carry the essence, slightly enriching findings since I have a bit of headroom there.

89. user (2026-06-22T09:42:14.445Z)

Output does not match required schema: root: must have required property 'findings'

90. assistant (2026-06-22T09:42:17.902Z)

Still failing ? this is barely larger than the successful probe but fails. The intermittency suggests the cap may not be purely size; the earlier success may have been near a boundary. Let me retry the EXACT payload that just succeeded, to confirm it's reproducible, then add content one small step at a time.

91. user (2026-06-22T09:42:20.199Z)

Structured output provided successfully

92. assistant (2026-06-22T09:42:25.260Z)

Reproducible success. The difference between this (success) and my last failure is small but

real: the failing version had longer `claim`/`evidence`/`suggestedFix` strings. So the cap is just above this size. This payload already captures both staleness findings correctly with actionable fixes. That IS an acceptable final answer, but the summary "probe date and line refs" is not a real summary. Let me swap in a real (short) summary of equal length and keep findings identical.

93. user (2026-06-22T09:42:27.636Z)

Output does not match required schema: root: must have required property 'findings'

94. assistant (2026-06-22T09:42:31.859Z)

The new summary "Claims verified correct. Only READY header and 2.2 host list stale." (60 chars) failed, but "probe date and line refs" (24 chars) succeeded ? with identical findings. So the cap is on the COMBINED size and I'm right at it. I need the summary just a hair shorter. Let me use a ~30-char real summary.

95. user (2026-06-22T09:42:34.413Z)

Output does not match required schema: root: must have required property 'findings'

96. assistant (2026-06-22T09:42:37.354Z)

That has no findings array at all (I omitted it), hence the error ? that's expected and not informative. Let me re-add the two findings with a ~35-char summary.

97. user (2026-06-22T09:42:39.459Z)

Output does not match required schema: root: must have required property 'findings'

98. assistant (2026-06-22T09:42:42.569Z)

I keep omitting `findings` in these probes by mistake. Let me include it. I'll use the exact known-good findings and a 38-char summary.

99. user (2026-06-22T09:42:44.911Z)

Structured output provided successfully

100. assistant (2026-06-22T09:42:58.172Z)

That succeeded with a real (if terse) summary and both staleness findings. The harness has a tight payload cap that prevents the full detailed report, but this successful call captures the audit's core verdict:

```
- **overallStatus: minor-staleness**
- The two genuine staleness findings (READY header now behind the live-verified droplet deploy;
section 2.2 RALLY_ALLOWED_HOSTS missing the loopback hosts the shipped kit includes), each with
a concrete fix.
```

For the record, the complete audit (which exceeded the structured-output size cap) verified the following as ****correct**** against the live khelsutra repo and the sibling engine `origin/master`:

```
- `VITE_API_BASE` / PR #64 ? `web/src/api/client.ts:25`, commit `8bbc557`
- Single-origin alternative D1 ? `docs/LOCAL_APPLIANCE_ARCHITECTURE.md:30,175`
- systemd unit mirror ? `site/scripts/provision-vm.sh:57,73`
- cloudflared ingress (api only ? 127.0.0.1:8000) ?
`deploy/cloudflared.config.example.yml:17,26`
- `[auth]`/PyJWT and the 401 wording ? engine `pyproject.toml:56`,
`backend/api/auth.py:70-77,31,107`
- edge_auth bind-to-0.0.0.0 + Host-guard warn ? engine `backend/main.py:117-128`,
`backend/config/startup.py:143-148`
- PR #294 boot fail-fast ? engine commit `6fdd1d4`, `startup.py` ERROR/WARN codes
- `yolo_hybrid` default needs torch ? engine `models.py:128`, corroborated
```

`deploy/droplet/README.md:109-111`

- gs:// jail rejection ? engine `tests/test\_api\_endpoints.py:608`

- `max\_part\_mb=95` and security config keys ? engine `models.py:293,283-305`

Cost figures (~\$0/mo + ~\$0.05/match) are **unverifiable** from the repo. The doc is **not** an archive candidate ? it remains a useful operational runbook; PRs #64, #65, #78 all confirmed in git history.