

claude-windows / 8d4866bf-8edd-4da9-909e-022ddee12675

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\8d4866bf-8edd-4da9-909e-022ddee12675\subagents\workflows\wf_3d035317-127\agent-a558218cfbd1b5ec4.jsonl`  
- SHA-256: `f30e1eb26e475f9a9c43847273797d7a67d4a5a9cb263700d7d31dacc1e79031`  
- Source modified: `2026-06-22T09:47:18+00:00`  
- Imported at: `2026-07-05T16:48:25+00:00`  
- project: `wf_3d035317-127`  
- session_id: `8d4866bf-8edd-4da9-909e-022ddee12675`
```

Transcript

1. user (2026-06-22T09:46:23.307Z)

You are fixing staleness in ONE doc of the KhelSutra repo, applying an audit that already ran. Be SURGICAL: minimal, targeted edits that match the audited fix; preserve the doc's voice, structure, and markdown formatting. Do NOT wholesale-rewrite a doc.

KhelSutra repo root: C:\Users\avidu\Projects\khelsutra
ENGINE repo root (sibling, for cross-repo verification): C:\Users\avidu\Projects\badminton-highlight-indexer (origin/master freshly fetched)
Today: 2026-06-22. KhelSutra is on branch docs/freshness-sweep (off master).

YOUR DOC (edit ONLY this file): C:\Users\avidu\Projects\khelsutra\docs\DRY_RUN_CUJ.md

The audit findings are in this JSON file: C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-indexer\8d4866bf-8edd-4da9-909e-022ddee12675\tasks\w3dbzrpby.output
Shape: { result: { perDoc: [{doc, overallStatus, findings, summary}], staleFindings: [{doc, claim, docLocation, verdict, evidence, suggestedFix, severity}] } }.
Read it and select ONLY entries where doc === "docs/DRY_RUN_CUJ.md". Your worklist = that doc's staleFindings. Also skim that doc's perDoc.findings for "unverifiable" items ? do NOT act on those unless you can now positively verify them.

For each stale finding:

1. RE-VERIFY it yourself against the LIVE repo before editing ? Read the doc lines + the cited code. For any claim about the ENGINE (paths like backend/main.py, /api/*, native_wasb_runner.py, auth/CORS), verify against the engine repo's origin/master: e.g. ``git -C /c/Users/avidu/Projects/badminton-highlight-indexer grep -n "<symbol>" origin/master -- backend/`` or ``git -C ... show origin/master:backend/main.py``. Trust nothing blindly ? the audit itself may be slightly off and engine line numbers DRIFT between PRs.
2. If confirmed, apply the suggestedFix as a SURGICAL Edit. For brittle ENGINE line-number anchors, prefer SYMBOL-based references (e.g. ``backend/main.py::get_config``) over chasing exact line numbers ? but only after confirming the symbol exists on engine origin/master.
3. DOCS ONLY: if a fix would require a CODE change (e.g. a comment in site/server.js), do NOT edit code ? record it under codeChangesNeeded.

Hard rules:

- Edit ONLY your assigned doc file. Run NO git WRITE commands (no add/commit/checkout/branch/stash) ? read-only git for verification only.
- Keep status markers HONEST: flip Todo->Done only WITH the PR number / file evidence; never invent a completion.
- If a finding is wrong, not reproducible, or too risky to auto-apply, SKIP it and explain under skipped.
- Do not "improve" beyond the findings; no scope creep.

Return the structured report.

2. assistant (2026-06-22T09:46:25.436Z)

I'll start by reading the audit findings and the target doc.

3. user (2026-06-22T09:46:27.808Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-591 of 1722 total (61859 tokens, cap 25000). Call Read with offset=592 limit=591 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-reminder>

```
1      {
2      "summary": "Audit every KhelSutra doc for staleness against the live repo, returning
structured findings",
3      "agentCount": 15,
4      "logs": [
5      "audited 15/15 docs; 44 stale findings; archive-candidates: none"
6      ],
7      "result": {
8      "perDoc": [
9      {
10     "doc": "NEXT_STEPS.md",
11     "overallStatus": "minor-staleness",
12     "summary": "NEXT_STEPS.md is overwhelmingly accurate against master: routes,
file paths, the Worker/D1 names, the 6-language localization, the 3-step wizard, the per-rally
PR numbers (#11/#14/#21/#28/#30/#31), the appliance/per-rally doc statuses, and every still-
pending \"Queued\" item all check out against code and git. The only real staleness is the \"PRs
through #93 merged\" count, which lags by one because PR #94 (the very commit that wrote
\"through #93\") merged afterward; plus a low-severity GCS-artifact filename mismatch. Not an
archive candidate ? it is the live next-steps tracker.",
13     "findings": [
14     {
15     "claim": "PRs through #93 merged (master).",
16     "docLocation": "NEXT_STEPS.md:13",
17     "verdict": "stale",
18     "evidence": "git log --oneline -1 = '50b2bd7 chore(repo): ... correct stale
NEXT_STEPS claims (#94)'; total 94 commits/PRs on master. Ironically #94 is the commit that
updated this line to '#93', then merged itself, so the doc now lags one PR behind master.",
19     "suggestedFix": "Update to 'PRs through #94 merged (master).'",
20     "severity": "medium"
21     },
22     {
23     "claim": "Where we are (2026-06-20): ... khelsutra.guru is LIVE and fully
shipped (landing + /capture checklist + OG image + segmented Request-Access form).",
24     "docLocation": "NEXT_STEPS.md:5-14",
25     "verdict": "correct",
26     "evidence": "site/server.js:98 maps /capture->capture.html;
site/public/og.png + capture.html tracked; index.html:191-193 textareas + lines 537-581
academy/individual/gpu/storage fields. Header date 2026-06-20 is slightly behind the latest work
(PR #94 merged 2026-06-22) but the substance holds.",
27     "suggestedFix": "Optionally bump the 'Where we are' date to 2026-06-22 to
match the latest merged work.",
28     "severity": "low"
29     },
30     {
31     "claim": "On Cloudflare (Worker hidden-king-khelsutra + D1 khelsutra-
access): the form persists to D1 and emails a summary via an Email Routing send binding.",
32     "docLocation": "NEXT_STEPS.md:7-9",
33     "verdict": "correct",
34     "evidence": "site/wrangler.toml: name = \"hidden-king-khelsutra\";
[[d1_databases]] database_name = \"khelsutra-access\"; [[send_email]] name = \"NOTIFY\".
site/src/index.js:80 handles POST /api/request-access -> insert into D1 binding DB.",
35     "suggestedFix": "",
36     "severity": "low"
37     },
38     ]
39     },
40     }
```

```

38         {
39             "claim": "CI gates every PR ? three jobs: full site npm test, web/ SPA build
+ coverage floor, and a real-browser layout check.",
40             "docLocation": "NEXT_STEPS.md:9-12",
41             "verdict": "correct",
42             "evidence": "site/package.json scripts include test (node --test),
test:layout (layout.browser.mjs); web has a coverage floor (#80). PR #94 message confirms 'three
jobs: site npm test = 86 green, web SPA build+coverage, real-browser layout'. Non-browser site
suite ran green (61 tests in the 4 non-layout files).",
43             "suggestedFix": "",
44             "severity": "low"
45         },
46         {
47             "claim": "The form also runs on a dependency-free Node + SQLite host
(site/server.js), live on the DO droplet -> http://168.144.126.176:8787 (systemd khelsutra-
site).",
48             "docLocation": "NEXT_STEPS.md:11-13",
49             "verdict": "correct",
50             "evidence": "site/server.js exists, uses node:sqlite (header comment line
2). IP 168.144.126.176 + port 8787 + 'khelsutra-site' all appear in
docs/LOCAL_APPLIANCE_ARCHITECTURE.md:48 ('/srv/khelsutra-site, node :8787') and deploy refs.
Live-status is external/unverifiable from the repo, but the host code and identifiers are
correct.",
51             "suggestedFix": "",
52             "severity": "low"
53         },
54         {
55             "claim": "Docs: docs/REQUEST_ACCESS_FORM.md, docs/PORTABILITY.md.",
56             "docLocation": "NEXT_STEPS.md:13-14",
57             "verdict": "correct",
58             "evidence": "git ls-files: both docs/REQUEST_ACCESS_FORM.md and
docs/PORTABILITY.md are tracked.",
59             "suggestedFix": "",
60             "severity": "low"
61         },
62         {
63             "claim": "DONE ? Friendly & Multilingual (ARCHIVED 2026-06-22): localized to
all 6 languages (en/kn/hi/es/da/id); full record docs/archives/FRIENDLY_AND_MULTILINGUAL.md.",
64             "docLocation": "NEXT_STEPS.md:17-21",
65             "verdict": "correct",
66             "evidence": "docs/archives/FRIENDLY_AND_MULTILINGUAL.md is tracked; git log
shows PR #93 'docs: archive the Friendly & Multilingual tracker'. web/src/i18n/index.ts:25
SUPPORTED_LOCALES = [en,kn,hi,es,da,id]; PR #87 added es/da/id to the marketing site.",
67             "suggestedFix": "",
68             "severity": "low"
69         },
70         {
71             "claim": "Interactive per-rally selection MVP/M0 COMPLETE; Done: P0 design
(#11), P1 scaffold (#14), P2 grid+hook (#21), P3 parseQuery query bar (#28), P4 SelectionFooter
(#30), P5 real load+download (#31). Tracked doc docs/PER_RALLY_SELECTION.md (IN PROGRESS).",
72             "docLocation": "NEXT_STEPS.md:32-41",
73             "verdict": "correct",
74             "evidence": "git log: cb74ce5 #11 'tracked design ... RallyPicker'; aac7e0c
#14 'P1 scaffold'; 62598e7 #21 'P2 selectable rally grid + hook'; d537d19 #28 'P3 deterministic
parseQuery() query bar'; aff8de8 #30 'P4 SelectionFooter + min_shot_count warning'; 8bf200d #31
'P5 load real rallies + Build->compile->preview->download'. docs/PER_RALLY_SELECTION.md status
header = 'IN PROGRESS ? MVP/M0 (P1-P5) COMPLETE'.",
75             "suggestedFix": "",
76             "severity": "low"
77         },
78         {
79             "claim": "P2c first-run setup wizard ? 3-step SetupWizard from
/api/capabilities, 6 languages.",
80             "docLocation": "NEXT_STEPS.md (implied via track) and verified for the
merged #88",

```

```

81         "verdict": "correct",
82         "evidence": "web/src/components/SetupWizard.tsx:22 'const STEPS = 3'; line 1
comment 'Three quick steps'; line 3 'Localized (en/kn/hi/es/da/id)'; web/src/api/client.ts:139
GET /capabilities. git: 7ae9cee #88.",
83         "suggestedFix": "",
84         "severity": "low"
85     },
86     {
87         "claim": "Single-box appliance ? Tracked design
docs/LOCAL_APPLIANCE_ARCHITECTURE.md (DRAFT, owner-gated).",
88         "docLocation": "NEXT_STEPS.md:52-58",
89         "verdict": "correct",
90         "evidence": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md status header: 'Status:
DRAFT ? owner-gated (auth choice, retention, cost ceiling, GPU host, P5 bridge)'.",
91         "suggestedFix": "",
92         "severity": "low"
93     },
94     {
95         "claim": "Queued: Always-Use-HTTPS + HSTS (Cloudflare) and Cloudflare Web
Analytics token paste into site/public/index.html ? both still pending.",
96         "docLocation": "NEXT_STEPS.md:44-45",
97         "verdict": "correct",
98         "evidence": "No Strict-Transport-Security/HSTS header found in site/*.js or
wrangler.toml. site/public/index.html:629 still has data-cf-beacon token 'REPLACE_WITH_CF_TOKEN'
(commented-out block). Both correctly listed as not-yet-done.",
99         "suggestedFix": "",
100        "severity": "low"
101    },
102    {
103        "claim": "Droplet -> production: add VM-path SMTP so the Node host also
emails (currently noopNotify).",
104        "docLocation": "NEXT_STEPS.md:46-47",
105        "verdict": "correct",
106        "evidence": "site/server.js:14 imports noopNotify; line 24 'const notify =
noopNotify(); // VM host: no email yet (add SMTP later)'.",
107        "suggestedFix": "",
108        "severity": "low"
109    },
110    {
111        "claim": "Org consolidation DONE; khelsutra (this site repo) transfer is an
owner step still PENDING (unchecked).",
112        "docLocation": "NEXT_STEPS.md:65-73",
113        "verdict": "correct",
114        "evidence": "git remote -v shows origin =
https://github.com/avidullu/khelsutra.git (NOT under Khelsutra org) ? consistent with the '?
owner step pending' marker. Other repos already moved: docs/DRY_RUN_CUJ.md:49 links
github.com/Khelsutra/badminton-highlight-indexer#289.",
115        "suggestedFix": "",
116        "severity": "low"
117    },
118    {
119        "claim": "A real-UI CUJ screencast exists -> promo_cuj.mp4/.gif in
gs://khelsutra-rally-corpus/promo/.",
120        "docLocation": "NEXT_STEPS.md:29-31",
121        "verdict": "unverifiable",
122        "evidence": "No GCS access to confirm the object. The companion doc
docs/DRY_RUN_CUJ.md:85 names the actual artifacts 'khelsutra-bucket-cuj-screencast.mp4' and ':60
khelsutra-i18n-screencast.mp4' in the same gs://khelsutra-rally-corpus/promo/ prefix;
'promo_cuj.mp4' appears nowhere else in the repo, suggesting a possible filename mismatch.",
123        "suggestedFix": "Reconcile the filename with docs/DRY_RUN_CUJ.md (likely
'khelsutra-bucket-cuj-screencast.mp4') or confirm 'promo_cuj.mp4' is a real object in the
bucket.",
124        "severity": "low"
125    },
126    {

```

```

127         "claim": "Control plane (control-plane/) at beta.",
128         "docLocation": "NEXT_STEPS.md:62",
129         "verdict": "correct",
130         "evidence": "control-plane/ directory does not exist in the repo (test -d
looking/aspirational item, not a claim of present existence ? so it is not stale as written.",
131         "suggestedFix": "",
132         "severity": "low"
133     }
134 ]
135 },
136 {
137     "doc": "README.md",
138     "overallStatus": "significant-staleness",
139     "summary": "The repo-structure list, deploy references, and canonical-doc
pointers are all correct, but the README's headline status is badly out of date: it still calls
the repo a \"pre-alpha scaffold\" whose SPA has yet to land \"per milestone M1\", whereas the
web/ SPA is in fact built, multilingual, and shipped (per-rally MVP P0?P5, P2c wizard, in-UI
upload, 6 languages; khelsutra.guru is LIVE). The control-plane/ directory it lists also does
not yet exist.",
140     "findings": [
141     {
142         "claim": "Status: pre-alpha scaffold. SPA lands in `web/` per execution-plan
milestone M1.",
143         "docLocation": "README.md:9",
144         "verdict": "stale",
145         "evidence": "web/src/ has 51 tracked source files incl. web/src/App.tsx;
`git log --oneline -- web/` = 31 commits; merged PRs show the SPA is feature-complete: P2c
first-run wizard (#88), in-UI file upload T4 (#66), per-rally selection MVP P0?P5
(#11/#14/#21/#28/#30/#31), 6-language catalogs (#67). NEXT_STEPS.md:7 states 'khelsutra.guru is
LIVE and fully shipped'. UI_ALPHA_EXECUTION_PLAN.md:36 defines M1 only as the initial
'Vite+React+TS+Tailwind scaffold', which is long past.",
146         "suggestedFix": "Replace with a current status, e.g.: 'Status: alpha ? the
web/ SPA is built and multilingual (6 languages), with the per-rally selection MVP and first-run
setup wizard shipped; khelsutra.guru / app.khelsutra.guru are live (Cloudflare Pages +
cloudflared tunnel + Access). See NEXT_STEPS.md.'",
147         "severity": "high"
148     },
149     {
150         "claim": "the web app (alpha SPA) now, the hosted control plane at beta",
151         "docLocation": "README.md:4",
152         "verdict": "stale",
153         "evidence": "This tagline pairs with line 9's 'pre-alpha scaffold', which is
contradicted by the shipped, live SPA (see web/ evidence above and NEXT_STEPS.md:7). The 'alpha
SPA now' framing is directionally right but undercut by the stale 'pre-alpha scaffold' status it
sits beside.",
154         "suggestedFix": "Keep 'web app (alpha SPA) now' but align the rest of the
intro/status so it no longer describes the SPA as not-yet-landed.",
155         "severity": "medium"
156     },
157     {
158         "claim": "`control-plane/` ? (beta) Cloud Run FastAPI control plane",
159         "docLocation": "README.md:17",
160         "verdict": "stale",
161         "evidence": "No control-plane/ directory exists: `ls control-plane` -> 'No
such file or directory'; `git ls-files control-plane/` returns nothing. Top-level tracked dirs
are deploy/, docs/, locales/, site/, web/ only.",
162         "suggestedFix": "Mark it as a future/planned directory that does not yet
exist (e.g. move it out of the 'Structure' list into a 'Planned (beta)' note), so a reader
doesn't expect the path to be present.",
163         "severity": "medium"
164     },
165     {
166         "claim": "Engine ... sibling repo `badminton-highlight-indexer` (rename
pending).",

```

```

167         "docLocation": "README.md:6",
168         "verdict": "correct",
169         "evidence": "Sibling repo still exists under that name: `ls
C:/Users/avidu/Projects/` shows 'badminton-highlight-indexer'; no rename has occurred, so
'rename pending' holds.",
170         "suggestedFix": "",
171         "severity": "low"
172     },
173     {
174         "claim": "Canonical planning docs live in the engine repo's docs/ for now:
UI_PROPOSAL.md, HOSTING_PLAN.md, UI_ALPHA_EXECUTION_PLAN.md.",
175         "docLocation": "README.md:7-8",
176         "verdict": "correct",
177         "evidence": "All three exist in the engine repo (badminton-highlight-
indexer/docs/UI_PROPOSAL.md, HOSTING_PLAN.md, UI_ALPHA_EXECUTION_PLAN.md present) and are NOT
tracked in khelsutra/docs (`git ls-files docs/` in khelsutra returns none of them; `ls
docs/UI_PROPOSAL.md` -> not found).",
178         "suggestedFix": "",
179         "severity": "low"
180     },
181     {
182         "claim": "`deploy/` ... alpha deploy kit ... DEPLOY_ALPHA.md ... +
cloudflared.config.example.yml",
183         "docLocation": "README.md:15",
184         "verdict": "correct",
185         "evidence": "git ls-files deploy/ shows deploy/DEPLOY_ALPHA.md and
deploy/cloudflared.config.example.yml (plus deploy/droplet/*). DEPLOY_ALPHA.md:1-5 matches the
described 'shareable page, no SSH; SPA on Pages + engine on api.* tunnel + CF Access' content.",
186         "suggestedFix": "",
187         "severity": "low"
188     },
189     {
190         "claim": "`site/` ? public marketing landing page (static; served at
khelsutra.guru)",
191         "docLocation": "README.md:13",
192         "verdict": "correct",
193         "evidence": "site/ exists with site/public/index.html, server.js, etc.;
DEPLOY_ALPHA.md and NEXT_STEPS.md confirm the marketing site is served at khelsutra.guru (apex)
while the SPA is at app.khelsutra.guru.",
194         "suggestedFix": "",
195         "severity": "low"
196     },
197     {
198         "claim": "`web/` ? React SPA (alpha; Vite + React + TS + Tailwind;
Cloudflare Pages build root)",
199         "docLocation": "README.md:14",
200         "verdict": "correct",
201         "evidence": "web/ is a Vite+React+TS+Tailwind SPA (web/README.md:3 'Vite +
React + TS + Tailwind v4'); DEPLOY_ALPHA.md:14 confirms 'app.khelsutra.guru -> Cloudflare Pages
(the web/ SPA)'. The tech-stack/build-root facts are accurate; only the inline 'alpha' maturity
ties into the stale status above.",
202         "suggestedFix": "",
203         "severity": "low"
204     },
205     {
206         "claim": "License: All rights reserved (private; no OSS license).",
207         "docLocation": "README.md:19-21",
208         "verdict": "correct",
209         "evidence": "No LICENSE file is tracked (not present in `git ls-files`);
consistent with a private 'all rights reserved' repo.",
210         "suggestedFix": "",
211         "severity": "low"
212     }
213 ]
214 },

```

```

215     {
216         "doc": "deploy/DEPLOY_ALPHA.md",
217         "overallStatus": "minor-staleness",
218         "summary": "Verified correct except a stale header",
219         "findings": [
220             {
221                 "claim": "header stale deploy live",
222                 "docLocation": "deploy/DEPLOY_ALPHA.md:3",
223                 "verdict": "stale",
224                 "evidence": "droplet/README.md:100 live-verified 2026-06-22",
225                 "suggestedFix": "set header to LIVE",
226                 "severity": "medium"
227             },
228             {
229                 "claim": "host list too narrow",
230                 "docLocation": "deploy/DEPLOY_ALPHA.md:66",
231                 "verdict": "stale",
232                 "evidence": "engine.env.example adds loopback",
233                 "suggestedFix": "add localhost",
234                 "severity": "low"
235             }
236         ]
237     },
238     {
239         "doc": "deploy/droplet/README.md",
240         "overallStatus": "fresh",
241         "summary": "This command-dense droplet quickstart is accurate and matches the
live repo: every cited kit file exists with the described contents, and the engine-behavior
claims (fail-fast boot posture, RALLY_* exposure vars, the Gemini-default-on-CPU gotcha, the
gs:// jail rejection, the --server entryptpoint) all verify against the engine repo and
DEPLOY_ALPHA.md. The only items not locally reproducible are the 2026-06-22 droplet live-run
figures (e.g. \"1313 passed\"), which are an off-box assertion ? they are internally plausible
(engine has ~1269 test functions before parametrization) but cannot be confirmed from these
repos.",
242         "findings": [
243             {
244                 "claim": "Four kit files exist: config.json.example, engine.env.example,
khelsutra-engine.service (and ../cloudflared.config.example.yml), each going to the documented
destination with the documented purpose.",
245                 "docLocation": "deploy/droplet/README.md:9-14",
246                 "verdict": "correct",
247                 "evidence": "git ls-files deploy/ lists all four:
deploy/droplet/{config.json.example,engine.env.example,khelsutra-engine.service} and
deploy/cloudflared.config.example.yml; contents read and match (config.json.example has the
security block + cf_team/aud; engine.env.example has the two RALLY_* exposure vars +
RALLY_APP_HOME; the service binds 127.0.0.1:8000; cloudflared.config routes api. ->
127.0.0.1:8000).",
248                 "suggestedFix": "",
249                 "severity": "low"
250             },
251             {
252                 "claim": "The boot fails fast (exit 1, '[STARTUP] refusing to start') if
edge_auth_enabled=true but the jail or cf_team_domain/cf_access_aud is unset.",
253                 "docLocation": "deploy/droplet/README.md:53-54",
254                 "verdict": "correct",
255                 "evidence": "backend/main.py:857 logs '[STARTUP] refusing to start';
backend/config/startup.py:109 EXPOSED_NO_VIDEO_ROOT and :123 EXPOSED_EDGE_AUTH_INCOMPLETE are
both LEVEL_ERROR (startup.py:41,272-276), while EXPOSED_AUTH_EXTRA_MISSING (PyJWT) is LEVEL_WARN
(:135). From engine PR #294 (commit 0c0e064).",
256                 "suggestedFix": "",
257                 "severity": "low"
258             },
259             {
260                 "claim": "indexing.default_segmenter MUST be gemini on a CPU box; the
shipped default yolo_hybrid needs torch/GPU and every job fails 'No module named torch' without

```

```

it.",
261         "docLocation": "deploy/droplet/README.md:109-113",
262         "verdict": "correct",
263         "evidence": "config.json.example:9-11 sets default_segmenter=gemini. Engine
AppConfig default is yolo_hybrid (backend/config/models.py:128), so an all-defaults load on the
droplet (which hand-writes config.json rather than running the --app LIGHT seeder) does land on
yolo_hybrid. The LIGHT 'motion' default (loader.py:227) only applies via
write_light_default_config (indexer --app first-run, PR #310), not this systemd path ? so the
gotcha is accurate for this deploy.",
264         "suggestedFix": "",
265         "severity": "low"
266     },
267     {
268         "claim": "RALLY_ALLOWED_HOSTS is the load-bearing fix; with edge_auth on,
the Host-header guard defaults loopback-only and 400s every tunneled request until set.",
269         "docLocation": "deploy/droplet/README.md:92, engine.env.example:14-17",
270         "verdict": "correct",
271         "evidence": "backend/config/startup.py:143-147 raises
EXPOSED_HOST_ALLOWLIST_LOOPBACK (LEVEL_WARN) when edge_auth is ON and RALLY_ALLOWED_HOSTS is
unset, with the exact 'will 400 forwarded requests' rationale; backend/main.py reads
RALLY_ALLOWED_HOSTS / RALLY_BIND_HOST / RALLY_CORS_ALLOW_ORIGINS. Matches DEPLOY_ALPHA.md §2.2
repro (400->200).",
272         "suggestedFix": "",
273         "severity": "low"
274     },
275     {
276         "claim": "Pasting a gs:// URI is rejected by the jail on the exposed box;
upload is the alpha path.",
277         "docLocation": "deploy/droplet/README.md:96-98",
278         "verdict": "correct",
279         "evidence": "tests/test_api_endpoints.py:608
test_process_hosted_mode_rejects_nonlocal_uri pins the gs:// -> 400 behavior under
allowed_video_root (referenced at :541 as the counterweight); the jail is required to expose
(config.json.example sets allowed_video_root).",
280         "suggestedFix": "",
281         "severity": "low"
282     },
283     {
284         "claim": "Pages build: Root dir web, Build npm run build, Output web/dist,
with VITE_API_BASE = https://api.khelsutra.guru/api.",
285         "docLocation": "deploy/droplet/README.md:73-76",
286         "verdict": "correct",
287         "evidence": "khelsutra web/ SPA exists (git ls-files web/...);
web/package.json build script = 'vite build'; web/src/api/client.ts:25 reads
import.meta.env.VITE_API_BASE (default '/api'), and client.test.ts:32 exercises the exact split
value https://api.khelsutra.guru/api. Vite default outDir 'dist' relative to root web/ =
web/dist.",
288         "suggestedFix": "",
289         "severity": "low"
290     },
291     {
292         "claim": "The shipped engine is run via the auth extra and python -m
backend.main --server, binding 127.0.0.1:8000.",
293         "docLocation": "deploy/droplet/README.md:33,47-49; khelsutra-
engine.service:26",
294         "verdict": "correct",
295         "evidence": "backend/main.py:867 defines the --server flag and :1107 calls
uvicorn.run(app, host=bind_host, port=port); --setup flag at :865; the service ExecStart uses
python -m backend.main --server.",
296         "suggestedFix": "",
297         "severity": "low"
298     },
299     {
300         "claim": "Live-verified on the droplet (2026-06-22): full suite 1313 passed
+ ruff/mypy clean on the box; boot-posture, Host-guard 400->200, path-jail, T4 upload->jail,

```



```

edge-auth 401-gate exercised live; full DNS->Cloudflare->Access->tunnel->engine chain
confirmed.",
301         "docLocation": "deploy/droplet/README.md:100-106",
302         "verdict": "unverifiable",
303         "evidence": "This is an off-box live-run assertion that cannot be reproduced
from the khelsutra or engine repos here. Sanity check: the engine tests/ dir has ~1269 'def
test_' functions (parametrize expansions push the collected/passed count higher), so '1313
passed' is plausible and consistent, not contradicted. Marked unverifiable per the file:line-
evidence discipline; no sign of staleness.",
304         "suggestedFix": "",
305         "severity": "low"
306     },
307     {
308         "claim": "Cost: Cloudflare Pages + Tunnel + Access are free (Access <= 50
users); standing cost ~ $0 + Gemini per run.",
309         "docLocation": "deploy/droplet/README.md:16-18",
310         "verdict": "unverifiable",
311         "evidence": "External Cloudflare pricing claim; not checkable against the
repo. Consistent with DEPLOY_ALPHA.md:3 ('~$0/mo standing ... + Gemini per-run'). No internal
contradiction.",
312         "suggestedFix": "",
313         "severity": "low"
314     }
315 ]
316 },
317 {
318     "doc": "docs/ANNOTATE_AND_TRAIN.md",
319     "overallStatus": "significant-staleness",
320     "summary": "A DRAFT tracker whose marketing-page deliverables (P0-P4 plus the
real annotator screenshots P5 mentions) are all shipped and merged (#15/#22/#23/#25), yet the
doc was last touched 2026-06-20 (#23) and never updated to reflect that. The core
engine/annotator technical claims (schema bridge, ship-gate refusing a commercial ship, MIT
annotator v1.6.4, guardrails) all still verify as correct ? only their file:line citations have
drifted. Main staleness: status still DRAFT despite shipped work, P5 marked Todo though
annotator screenshots are committed, the doc is NOT actually indexed in docs/README.md (claimed
three times), and the page since grew 6-language localization the tracker never anticipated. It
is an archive candidate once honestly closed out.",
321     "findings": [
322     {
323         "claim": "Status: DRAFT ? owner-gated on scope (header, line 3); Last
updated: 2026-06-20.",
324         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:3",
325         "verdict": "stale",
326         "evidence": "git log --follow shows the marketing-page work merged: P0 doc
(#15, 6b62ebf), own-model page P1-P4 (#22, b1110e1), org link (#23, 87520b8). Per $7, P0-P4 are
all marked Done ? yet the header still reads DRAFT. Lifecycle (line 4) is DRAFT->IN
PROGRESS->DONE; with P0-P4 shipped this is at least IN PROGRESS.",
327         "suggestedFix": "Flip Status to IN PROGRESS (only owner-gated P5-P9 remain)
and bump Last updated; or, if the owner considers the marketing surface complete, mark DONE and
archive per the lifecycle note.",
328         "severity": "medium"
329     },
330     {
331         "claim": "P5 'Capture + commit real screenshots ... replace placeholders' ?
Status ? (Todo), Gated Owner.",
332         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:76",
333         "verdict": "stale",
334         "evidence": "PR #25 (1906c82) 'show the real annotator screenshots'
committed site/public/img/annotate-enable.png (133712 bytes) and annotate-tool.png (153133
bytes); own-model.html:111-112 references them as the real VLC-UI shots (no people/footage). The
annotator-UI half of P5 is done and merged, not Todo.",
335         "suggestedFix": "Mark P5 ? for the annotator-UI screenshots (cite #25); if
footage screenshots with consented people are still pending, split that into its own row so the
done part isn't hidden under a blanket Todo.",
336         "severity": "high"

```

```

337         },
338         {
339             "claim": "Doc is 'index[ed] in docs/README.md' (line 5), §7 P0 'indexed in
docs/README.md' (? , #15), and DoD 'doc indexed in docs/README.md'.",
340             "docLocation": "docs/ANNOTATE_AND_TRAIN.md:5,71,91",
341             "verdict": "stale",
342             "evidence": "grep -ni 'annotate|own.model|train.your|ANNOTATE_AND'
docs/README.md returns NOT FOUND; the docs/README.md 'Tracked work' list (lines 1-30) does not
mention ANNOTATE_AND_TRAIN.md or own-model. The claimed index entry does not exist.",
343             "suggestedFix": "Either add an ANNOTATE_AND_TRAIN.md entry to
docs/README.md's 'Tracked work' list, or correct P0/DoD to reflect that it is not yet indexed.",
344             "severity": "high"
345         },
346         {
347             "claim": "Engine ship-gate (gate_verdict, harness.py:98) currently refuses a
commercial ship regardless of F1: WASB provenance C3 unresolved + F1@tIoU not calibrated
(#172).",
348             "docLocation": "docs/ANNOTATE_AND_TRAIN.md:12,31",
349             "verdict": "correct",
350             "evidence": "training/gen0/harness.py:177 returns {'ship': False, ...}; :175
'WASB checkpoint provenance C3 UNRESOLVED ? commercial ship blocked regardless of F1'; :174
'ship-gate F1@tIoU target NOT YET CALIBRATED ... see issue #172'. Fact correct; the cited line
(:98) has drifted ? gate_verdict is now at :142.",
351             "suggestedFix": "Update the line citations: gate_verdict is at
harness.py:142 (not :98).",
352             "severity": "low"
353         },
354         {
355             "claim": "Default synth path: synth_trajectory_from_labels uses
csv.DictReader reading row['start_time']/row['end_time'] (stub_runner.py:139) ? MATCHES the
annotator header exactly.",
356             "docLocation": "docs/ANNOTATE_AND_TRAIN.md:46",
357             "verdict": "correct",
358             "evidence": "backend/pipeline/detectors/stub_runner.py:124 def
synth_trajectory_from_labels; :137 for row in _csv.DictReader(f); :139
rallies.append((float(row['start_time']), float(row['end_time']))). Annotator header
rally_number,start_time,end_time,... confirmed at rally-annotator vlc/rally_annotator.lua:128
and docs/CSV_FORMAT.md:6. Claim accurate.",
359             "suggestedFix": "",
360             "severity": "low"
361         },
362         {
363             "claim": "Canonical GT reader load_gt_intervals accepts start/end OR
start_time/end_time (gt_loader.py:14-90).",
364             "docLocation": "docs/ANNOTATE_AND_TRAIN.md:47",
365             "verdict": "correct",
366             "evidence": "backend/eval/gt_loader.py:30 _START_NAMES = ('start',
'start_time'); :44 def load_gt_intervals; :70 headerless positional col0,col1 path. Both aliases
accepted as claimed.",
367             "suggestedFix": "",
368             "severity": "low"
369         },
370         {
371             "claim": "Bundle attests no_paid_llm_training_labels:True (harness.py:192)
and label source is human (harness.py:178).",
372             "docLocation": "docs/ANNOTATE_AND_TRAIN.md:33,60",
373             "verdict": "correct",
374             "evidence": "training/gen0/harness.py:236 'no_paid_llm_training_labels':
True; :222 'label_source': 'human (rally-annotator golden CSVs)'. Facts correct; cited lines
(:192, :178) have drifted to :236 and :222.",
375             "suggestedFix": "Update line citations: no_paid_llm_training_labels at
harness.py:236, label_source at :222.",
376             "severity": "low"
377         },
378         {

```

```

379         "claim": "cmd_train (__main__.py:231) pulls labels/, joins video<->label by
file MD5 + stem-match labels/<name>.rallies.csv <-> videos/<name>.<ext> (__main__.py:271-280).",
380         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:31,49",
381         "verdict": "correct",
382         "evidence": "backend/worker/__main__.py:332 def cmd_train; :359 lists
{corpus}/labels/; :383-384 fname.replace('.rallies.csv','') stem join. Logic present and
correct; cmd_train line drifted from :231 to :332.",
383         "suggestedFix": "Update cmd_train citation to __main__.py:332.",
384         "severity": "low"
385     },
386     {
387         "claim": "Reuse map: add a /own-model alias in site/server.js:49-51; Worker
env.ASSETS.fetch needs no change (site/src/index.js:84).",
388         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:52,73",
389         "verdict": "correct",
390         "evidence": "site/server.js:99 'pathname === \"/own-model\" ? \"/own-
model.html\"' and :41 LOCALIZED_PAGES includes own-model.html (alias shifted off the cited
:49-51). site/src/index.js:84 'return env.ASSETS.fetch(request)' with no /own-model branch ?
confirms the Worker needs no change.",
391         "suggestedFix": "Update the server.js alias citation (now ~:99, not
:49-51).",
392         "severity": "low"
393     },
394     {
395         "claim": "Reuse precedent: screenshots under site/public/img/ (precedent
site/og.png + make_og.py).",
396         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:52",
397         "verdict": "correct",
398         "evidence": "git ls-files: site/make_og.py and site/public/og.png both
exist; site/public/img/ holds annotate-enable.png, annotate-tool.png, studio-*.png. The og asset
path is site/public/og.png (the doc wrote 'site/og.png' loosely), make_og.py is at
site/make_og.py.",
399         "suggestedFix": "Minor: the OG image is at site/public/og.png (not
site/og.png).",
400         "severity": "low"
401     },
402     {
403         "claim": "SPA today: web/src/App.tsx is one static screen, no router
(web/package.json); a correction screen needs react-router-dom.",
404         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:53,64",
405         "verdict": "stale",
406         "evidence": "grep on web/package.json: no react-router dep (claim 'no
router' still TRUE). But App.tsx is no longer 'one static screen' ? it now imports and
conditionally renders SetupWizard (web/src/App.tsx:11,65,242, P2c first-run wizard from PR #88).
The 'one static screen' framing predates the wizard.",
407         "suggestedFix": "Reword to 'no router; App.tsx renders a first-run
SetupWizard + the ingest panel' so the SPA surface description matches the post-#88 code.",
408         "severity": "low"
409     },
410     {
411         "claim": "SPA correction route depends on a new, unbuilt engine endpoint
PATCH /api/segments/{id} (dead columns
candidate_segments.human_label/notes/confirmed_segment_id); P8 ? Gated.",
412         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:53,79",
413         "verdict": "correct",
414         "evidence": "grep across backend/ for 'api/segments'/'PATCH' on
candidate_segments finds only the candidate_segments table + query_candidate_segments readers
(database.py:217, eval/*.py); no PATCH /api/segments route exists. P8 is still genuinely unbuilt
as stated.",
415         "suggestedFix": "",
416         "severity": "low"
417     },
418     {
419         "claim": "Relation note: peer of LOCAL_APPLIANCE_ARCHITECTURE.md (the latter
lands with PR #13).",

```

```

420         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:6,84",
421         "verdict": "stale",
422         "evidence": "git ls-files confirms docs/LOCAL_APPLIANCE_ARCHITECTURE.md
already exists in the repo; it has long since landed, so the forward-looking 'lands with PR #13'
/ 'cross-link once PR #13 merges' phrasing is outdated.",
423         "suggestedFix": "Drop the 'lands with PR #13' future-tense framing ? the doc
exists now; add the cross-link.",
424         "severity": "low"
425     },
426     {
427         "claim": "rally-annotator is a single-file VLC Lua extension, VERSION 1.6.4,
MIT-licensed, public at github.com/Khelsutra/rally-annotator.",
428         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:29",
429         "verdict": "correct",
430         "evidence": "rally-annotator/vlc/rally_annotator.lua:1 '(v1.6.4)'; LICENSE
first line 'MIT License'; own-model.html:109/191 links to https://github.com/Khelsutra/rally-
annotator (org-link migration was PR #23, captured in the doc's own history).",
431         "suggestedFix": "",
432         "severity": "low"
433     },
434     {
435         "claim": "$7 tracker: P1/P2/P3/P4 all ? via '(this PR)'.",
436         "docLocation": "docs/ANNOTATE_AND_TRAIN.md:72-75",
437         "verdict": "correct",
438         "evidence": "own-model.html exists (P1); /own-model alias in server.js:99
(P2); nav cross-link index.html:255 + footer :620 (P3); verbatim annotator content + GitHub link
in own-model.html:108-191 (P4). All shipped ? git history shows these in PR #22 (b1110e1,
'annotate-train P1-P4'), so '(this PR)' should now read '#22'.",
439         "suggestedFix": "Replace the '(this PR)' placeholders in the PR column with
#22 (and #23 for the org-link follow-up) so the tracker cites real merged PRs.",
440         "severity": "medium"
441     }
442 ]
443 },
444 {
445     "doc": "docs/DRY_RUN_CUJ.md",
446     "overallStatus": "minor-staleness",
447     "summary": "The doc's substantive engineering claims are accurate and well-
grounded: the SPA back-half + B-P2 front-half ingest screen exist
(web/src/components/IngestPanel.tsx, web/src/hooks/useIngestJob.ts, #56), the engine bucket-
ingest routing is verified in code, and the three cited engine PRs (#280/#281/#289) plus all
four cited engine tests exist. The staleness is in metadata/framing: the header still says
Status: DRAFT though docs/README.md correctly calls it DELIVERED with DoD $6 boxes 1-3 done, and
the doc repeatedly describes the multilingual workflow as \"en/kn/hi\" (3 langs) when the SPA
now supports 6 (es/da/id added in #87). It is a completed tracker and an archive candidate once
the parked droplet box is closed.",
448     "findings": [
449         {
450             "claim": "Status: DRAFT ? owner-gated . ... Goal for this session.",
451             "docLocation": "docs/DRY_RUN_CUJ.md:3",
452             "verdict": "stale",
453             "evidence": "docs/README.md:4 marks this doc '? DELIVERED (2026-06-21)'; DoD
$6 boxes 1-3 are all checked [x] (DRY_RUN_CUJ.md:47-49) and the work shipped (#53/#55/#56/#57).
A DRAFT/'goal for this session' header no longer matches the completed, delivered state.",
454             "suggestedFix": "Flip the header to Status: DONE / DELIVERED (2026-06-21)
and drop 'Goal for this session'; note the only open box (droplet, $6 box 4) is owner-parked.
Consider archiving once that box closes.",
455             "severity": "medium"
456         },
457         {
458             "claim": "The SPA is 'Fully en/kn/hi (the switcher / ?lang=)' and the
multilingual workflow is en/kn/hi.",
459             "docLocation": "docs/DRY_RUN_CUJ.md:11",
460             "verdict": "stale",
461             "evidence": "web/src/i18n/index.ts:25 ? `export const SUPPORTED_LOCALES =

```

```

[\"en\\\", \"kn\\\", \"hi\\\", \"es\\\", \"da\\\", \"id\\\"]`; es/da/id were added to the SPA in PR #87
(cf59fbf 'add es/da/id to the marketing site (parity with the SPA's 6 languages)'). The SPA now
supports 6 languages, not 3.",
462     "suggestedFix": "Update 'en/kn/hi' to 'en/kn/hi/es/da/id (6 languages)'
where it describes the current switcher; the Hindi/Kannada run results in §7 stay accurate as
historical record.",
463     "severity": "low"
464 },
465 {
466     "claim": "This SPA front-half (ingest/process UI) is the *only* net-new work
for Path B (= B-P2); the SPA has no ingest/process UI today.",
467     "docLocation": "docs/DRY_RUN_CUJ.md:14",
468     "verdict": "correct",
469     "evidence": "B-P2 was net-new work at draft time and has since shipped:
web/src/components/IngestPanel.tsx + web/src/hooks/useIngestJob.ts (first landed in PR #56
'feat(web): B-P2 ? in-UI bucket ingest', 4525c5d, 2026-06-21). useIngestJob.ts drives POST
/api/process (processVideo) ? poll GET /api/jobs/{id} (getJob) ? hand video_id to App, exactly
as scoped.",
470     "suggestedFix": "",
471     "severity": "low"
472 },
473 {
474     "claim": "POST /api/process accepts gs://gcs://s3://local:// URIs and skips
the local-existence check for non-local backends; resolve via storage.resolve_input_ref;
gs://?gcs normalization.",
475     "docLocation": "docs/DRY_RUN_CUJ.md:15,68-72",
476     "verdict": "correct",
477     "evidence": "Engine backend/main.py:581 `input_ref =
storage.resolve_input_ref(req.video_path, ...)` and main.py:590 `if input_ref.backend ==
'local' and not os.path.exists(input_ref.key):` (the non-local skip), with the fetch at
main.py:458 `storage.acquire_local_input(...)`; backend/storage/ref.py:60 `if scheme == 'gs':
scheme = 'gcs'`. Substance matches; the exact line numbers 159/274/283 cited in the doc have
drifted as the engine shipped further PRs.",
478     "suggestedFix": "Optionally refresh the cited line numbers (now
~main.py:581/590/458, ref.py:60) since the engine surface moved; the routing logic itself is
intact.",
479     "severity": "low"
480 },
481 {
482     "claim": "Tested by
tests/test_api_endpoints.py::test_process_nonlocal_uri_skips_exists_and_fetches,
::test_process_hosted_mode_rejects_nonlocal_uri, ::test_process_unknown_scheme_returns_400,
::test_cli_nonlocal_input_fetches_and_keeps_uri.",
483     "docLocation": "docs/DRY_RUN_CUJ.md:15,76",
484     "verdict": "correct",
485     "evidence": "Engine repo grep: tests/test_api_endpoints.py:508
test_process_nonlocal_uri_skips_exists_and_fetches; :577
test_process_unknown_scheme_returns_400; :608 test_process_hosted_mode_rejects_nonlocal_uri;
:627 test_cli_nonlocal_input_fetches_and_keeps_uri. All four exist with the asserted names.",
486     "suggestedFix": "",
487     "severity": "low"
488 },
489 {
490     "claim": "B-P1 (gs://?/api/process?API-DB video_id) already shipped in
engine M2 / #280; the job-drain hang fixed in #281; the compile-from-bucket 500 fixed in #289.",
491     "docLocation": "docs/DRY_RUN_CUJ.md:15,63,68,88",
492     "verdict": "correct",
493     "evidence": "Engine git log: e111add 'feat(serving): M2 ? URI-aware input
... (#280)'; 41d45fd 'fix(server): drain async jobs under `python -m backend.main` (double-
module hang) (#281)'; e68b583 'fix(api): resolve a bucket-ingested video's source to a local
file before stitching (#289)'. All three PRs exist with descriptions matching the doc.",
494     "suggestedFix": "",
495     "severity": "low"
496 },
497 {

```

```

498         "claim": "DoD §6 box 3: in-UI Path-B bucket?reel CUJ DONE (live e2e
2026-06-21) ? 12 rallies loaded ? reel downloaded; box 4 (CPU droplet) still open/un-
provisioned.",
499         "docLocation": "docs/DRY_RUN_CUJ.md:49-50",
500         "verdict": "correct",
501         "evidence": "Box 3 is corroborated by PR #57 'docs(dry-run): Path-B live e2e
DONE ? in-UI bucket?reel CUJ + screencast (DoD §6 box 3 ?)' and the merged B-P2 SPA (#56).
docs/README.md:6-7 confirms box 4 (droplet, option 3) remains 'parked by the owner (GPU
subscription)'. The run-result specifics (12 rallies, reel sizes, GCS screencast paths) are
runtime artifacts not locally reproducible.",
502         "suggestedFix": "",
503         "severity": "low"
504     },
505     {
506         "claim": "Run-result specifics: 12 rallies / video_id 05e0e577 /
50.9MB+76.5s reels / screencasts at gs://khelsutra-rally-corpus/promo/*.mp4 and the staged
bucket video gs://khelsutra-rally-corpus/videos/mahadevpura_smoke_180s.mp4.",
507         "docLocation": "docs/DRY_RUN_CUJ.md:54-60,82-85",
508         "verdict": "unverifiable",
509         "evidence": "These are GCS bucket objects and a one-off local run's outputs;
they cannot be checked from the repo. No counter-evidence found ? the surrounding code/PR claims
that frame them are all verified.",
510         "suggestedFix": "",
511         "severity": "low"
512     }
513 ]
514 },
515 {
516     "doc": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md",
517     "overallStatus": "significant-staleness",
518     "summary": "A still-active DRAFT tracker (last updated 2026-06-21) whose §9
progress and several \"Absent (net-new)\" engine-fact anchors have drifted: the engine has since
shipped /api/capabilities, exposure-safety boot posture + an edge-auth layer, and merged
native_wasb_runner to master, while the alpha deploy went LIVE on 2026-06-22
(app.khelsutra.guru). The khelsutra-side claims (PR refs #11/#56/#64/#66, IngestPanel,
VITE_API_BASE) check out; the staleness is concentrated in engine status and the deploy/DoD
progress. Not yet an archive candidate (open owner gates remain) but needs a refresh pass.",
519     "findings": [
520     {
521         "claim": "§3 lists `GET /api/capabilities` under 'Absent (net-new): ...
`[verified ? grep returns none]`; §4.6 Reuse map lists it under 'Net-new'; §9 T7 (capabilities
+ health + banner) is ? Todo.",
522         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:49,113,164",
523         "verdict": "stale",
524         "evidence": "engine origin/master backend/main.py:340
`@app.get(\"/api/capabilities\")` (and main.py:307 `/api/health`, main.py:334 `/api/version`)
all exist on master. git grep '/api/capabilities' origin/master returns a hit.",
525         "suggestedFix": "Move `/api/capabilities` out of 'Absent (net-new)' and out
of §4.6 'Net-new'; mark the engine half of T7 (capabilities+health) as shipped (?/?), leaving
only the SPA banner pending.",
526         "severity": "high"
527     },
528     {
529         "claim": "TL;DR + Risk table: 'the engine has zero app-layer auth and CORS
`*` `[verified backend/main.py:42]` ? structural safety only because the engine binds
127.0.0.1.",
530         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:14,30,103,128",
531         "verdict": "stale",
532         "evidence": "engine origin/master backend/main.py:82 uses
`allow_origins=cors_origins_from_env()` (default ['*'] but env-lockable) with
`allow_credentials=False`; main.py:109 adds TrustedHostMiddleware (DNS-rebinding/Host
allowlist); backend/api/auth.py + config/startup.py enforce_startup_checks + M9 `Cf-Access-
Authenticated-User-Email` `owner_id` now exist on master. CORS is no longer a hardcoded bare
`*`, and an app-layer edge-auth integration now exists.",
533         "suggestedFix": "Reword to: CORS default is `*` but now env-configurable

```

```

(allow_credentials off), and the engine now ships a Host-header guard + an edge-auth integration
(Cf-Access trust) + exposure-safety boot posture ? so 'zero app-layer auth' is no longer
literally true. Re-anchor the line (main.py:42 no longer points at CORS).",
534         "severity": "high"
535     },
536     {
537         "claim": "$3 / $4.3: 'Native WASB is in-flight on branch `feat/native-wasb-
runner`: `native_wasb_runner.py` exists `[verified]`'.",
538         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:45,96",
539         "verdict": "stale",
540         "evidence": "engine origin/master now contains
backend/pipeline/detectors/native_wasb_runner.py and tests/test_native_wasb_runner.py (git ls-
tree origin/master). The runner has merged to master; the `feat/native-wasb-runner` branch still
exists but the file is no longer branch-only. (trajectory_hybrid._build_native_runner still
raises NotImplementedError at trajectory_hybrid.py:61 ? that sub-claim is still correct.)",
541         "suggestedFix": "Change 'in-flight on branch feat/native-wasb-runner' to
'native_wasb_runner.py is on master; trajectory_hybrid still refuses, fusion builds it; owner
byte-parity gate still pending.'",
542         "severity": "medium"
543     },
544     {
545         "claim": "$9 alpha-deploy block: 'the actual stand-up (tunnel + Pages + CF
Access + key rotation) is the **owner** step'; $11 DoD box 1 (CUJ-1 e2e on droplet behind gate)
and box 2 (raw :8000 refused, tested) unchecked.",
546         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:175,182,183",
547         "verdict": "stale",
548         "evidence": "deploy/droplet/README.md:100-106 '? Live-verified on the
droplet (2026-06-22)': engine booted on 168.144.126.176; suite 1313 passed; boot-posture fail-
fast, Host-guard 400?200, path-jail, T4 upload?jail, and edge-auth 401-gate all exercised live;
full DNS?Cloudflare?Access?tunnel?engine chain confirmed; app.khelsutra.guru up. Only Gemini-key
rotation + a browser OTP login remain.",
549         "suggestedFix": "Update $9 to note the alpha stand-up is LIVE (2026-06-22,
app.khelsutra.guru, posture test exercised) with only key rotation + browser-OTP login
outstanding; tick/qualify the corresponding DoD boxes.",
550         "severity": "high"
551     },
552     {
553         "claim": "Status header: 'Status: DRAFT ... Last updated: 2026-06-21' with
$9 the source of truth.",
554         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:3",
555         "verdict": "stale",
556         "evidence": "khelsutra git log shows merges through #94 (master HEAD
50b2bd7) and the alpha went live 2026-06-22, both after the doc's 2026-06-21 'last updated'.
Multiple $9 rows and engine facts have moved since (T7 engine half, alpha deploy, native_wasb on
master).",
557         "suggestedFix": "Bump 'Last updated' and consider promoting DRAFT?IN
PROGRESS; refresh $9 rows to current reality.",
558         "severity": "medium"
559     },
560     {
561         "claim": "D4 'Embed RallyPicker ... PR [#11]' and $9 T4 '[#66]', T12 'SPA
#56', alpha-deploy '[#64]' / 'VITE_API_BASE'.",
562         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:33,161,169,175",
563         "verdict": "correct",
564         "evidence": "khelsutra git log: cb74ce5 '...RallyPicker (#11)'; ca619fb 'T4
? in-UI file upload (#66)'; 4525c5d 'B-P2 ? in-UI bucket ingest (#56)'; 8bbc557 'VITE_API_BASE
(#64)'. web/src/components/IngestPanel.tsx handles both chunked upload (T4) and a gs:// URI
(T12); web/src/api/client.ts:25 reads VITE_API_BASE.",
565         "suggestedFix": "",
566         "severity": "low"
567     },
568     {
569         "claim": "T12 (bucket-URI ingest) marked ?: 'engine accepts URIs (M2/#280) +
SPA ingest screen shipped (B-P2, #56); in-UI gs:///reel CUJ proven live e2e (DRY_RUN_CUJ.md
$9)'.",

```

```

570         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:169",
571         "verdict": "correct",
572         "evidence": "docs/DRY_RUN_CUJ.md:80-92 §9 records the live in-UI gs://?reel
CUJ (12 rallies ? reel downloaded) and the compile-from-bucket 500 fix (engine #289);
IngestPanel.tsx accepts a gs:// URI. khelsutra #56 (B-P2) is merged.",
573         "suggestedFix": "",
574         "severity": "low"
575     },
576     {
577         "claim": "T11 retention marked ?: sweep tool shipped (engine #299), auto-
scheduling tracked (engine #301).",
578         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:168,130",
579         "verdict": "unverifiable",
580         "evidence": "These reference the engine repo's GitHub PR #299 / issue #301
(Khelsutra/badminton-highlight-indexer). Not falsifiable from the local khelsutra checkout; the
local engine checkout is on a feature branch and I did not confirm the specific sweep tool by PR
number. The khelsutra-side row metadata is internally consistent (#69 flipped T11 to ? per git
log bea8a2a).",
581         "suggestedFix": "",
582         "severity": "low"
583     },
584     {
585         "claim": "§3 Droplet reality: co-hosts the marketing site at
`/srv/khelsutra-site`, node `:8787`, engine at `~/rally`; MinIO 9000/9001 + GCS emulator 9023.",
586         "docLocation": "docs/LOCAL_APPLIANCE_ARCHITECTURE.md:48",
587         "verdict": "unverifiable",
588         "evidence": "These are live-droplet filesystem/port facts (168.144.126.176)
not checkable from the local repo. deploy/DEPLOY_ALPHA.md:31 and deploy/droplet/README.md:4
corroborate the IP/CPU-no-GPU/blr1 framing, but the exact paths/ports (/srv/khelsutra-site,
:8787, ~/rally) are not in-repo verifiable.",
589         "suggestedFix": "",
590         "severity": "low"
591     },

```

4. user (2026-06-22T09:46:27.840Z)

```

1      # Dry-run CUJ ? UI drives the tool on a bucket video ? reel downloaded local
(multilingual)
2
3      > **Status:** `DRAFT` ? owner-gated . **Owner:** Avi Dullu . **Created:** 2026-06-21 .
**Goal for this session.**
4      > **Umbrella:** a milestone of `docs/LOCAL_APPLIANCE_ARCHITECTURE.md` (the appliance) +
`docs/PER_RALLY_SELECTION.md` (the shipped SPA back-half).
5      > **Honesty note:** grounded in a file:line integration-surface map (this session);
claims tagged `[verified]` / `[gap]`.
6
7      ## 0. The goal (owner's words)
8      Operate the **web UI** ? point at a **video on the GCS bucket** ? the engine (tool)
processes it ? **review/select rallies in the UI** ? **compile** ? **download the reel to the
local machine**. Then **repeat with the UI in Hindi** to prove the multilingual workflow.
9
10     ## 1. What already works vs the gap `[verified]`
11     **? Back half (shipped ? PER_RALLY_SELECTION M0):** the SPA loads `?video=<id>` ? `POST
/api/query` (rallies, default all-selected) ? deterministic query/select ? `POST /api/compile` ?
`<video>` preview + **`<a download>` saves the MP4 locally**. Reel download-to-local **works
today** (`GET /api/highlights/{video_id}/{filename}` streams it; `/api/compile` also returns
`download_url`). Fully **en/ku/hi** (the switcher / `?lang=`).
12
13     **? Front half ? the gap is the SPA, NOT the engine `[corrected 2026-06-21; see §8]`:**
14     - The SPA has **no ingest/process UI** ? it can't point at a video (bucket or local),
can't call `/api/process`, can't poll `/api/jobs`. It only gets a `video_id` from `?video=` in
the URL. **This SPA front-half is the *only* net-new work for Path B (= B-P2).**
15     - **HTTP bucket-ingest ALREADY EXISTS in the engine (shipped M2 / [#280]) `[verified
file:line]`:** `POST /api/process` accepts a `gs:///`/`gcs:///` (also `s3:///`/`local:///`)
`video_path` directly. Submit resolves the URI and **skips the local-existence check for non-

```


local backends** (`backend/main.py:274,283` ? `storage.resolve\_input\_ref`; `gs://`?`gcs` normalized in `backend/storage/ref.py:60`); the job worker then fetches it to scratch (`\_execute\_processing` ? `storage.acquire\_local\_input`, `main.py:159`) and runs the **same** hash ? validate ? segment ? DB pipeline as a local file ? writing a queryable `video\_id` into the API's SQLite `videos`/`play\_segments`. Proven end-to-end by `tests/test\_api\_endpoints.py::test\_process\_nonlocal\_uri\_skips\_exists\_and\_fetches` (gs:// ? 202, not 404 ? job `done` ? `db.get\_video(video\_id)` row with the **source URI preserved**), plus the hosted-mode-rejection / unknown-scheme-400 / CLI-variant tests.

16 - **Where the earlier "net-new" claim went wrong:** it conflated the **API path** (`/api/process`, which DID gain URI ingest in M2) with the separate **worker CLI** (`python -m backend.worker infer --video-uri gcs://`), which writes `intervals.csv`/`report.json` to the bucket ? **not** the API DB. The CLI distinction is real; the "the endpoint is local-only" conclusion was not. ? **Engine-surface claims drift between sessions** (the engine ships PRs independently) ? re-verify any `[gap]`/`net-new` engine claim against the cited `file:line` before scoping it as work.

17
18 **## 2. Two paths to the goal**
19 ****Path A ? Cheapest dry run NOW (~zero new code).** Pre-process the bucket video out-of-band to mint a `video\_id` in the engine DB, then drive the existing SPA back-half. Proves the **multilingual UI-driven workflow** end-to-end. The only non-UI step is "point at the bucket" (a one-line CLI pre-process).

20
21 ****Path B ? Full "point at a bucket video IN the UI" (SPA-only ? smaller than first scoped).** The **HTTP bucket-ingest endpoint already exists** (§1) ? so Path B is **only the SPA front-half** (appliance S1/S2): an **ingest** input (enter a `gs://` URI) ? **POST /api/process** ? an **ingest + progress** screen polling **GET /api/jobs/{id}** ? on `done`, load **result.video_id** (which the existing back-half already renders). **No net-new engine endpoint** ? B-P1 is already done; the remaining work is **B-P2 (the SPA)**.

22
23 **## 3. Path A runbook (the cheapest dry run) `[verified-steps]`**
24 1. **Engine host + deps:** run the engine FastAPI (`python -m backend.main --server`, binds `127.0.0.1:8000`) on a box with `ffmpeg` + OpenCV. Gemini/motion paths are **CPU-only**, no WSL ? local Windows works; the CPU droplet works too. (GPU/WSL only needed for native WASB.)
25 2. **Get the bucket video local + into the API DB:**
26 `gsutil cp gs://<bucket>/<match>.mp4 ./incoming/` (or the worker fetch), then **POST /api/process** {"video_path": "<abs local path>", "segmenter": "gemini"} ? poll **GET /api/jobs/{job_id}** until **result.video_id**. (*`gemini` needs **GEMINI_API_KEY**; `motion` is key-free but hides fabricated metadata.)*
27 3. **Open the SPA against the engine:** `npm run dev` (Vite proxies `/api` ? `:8000`, same-origin) ? **http://localhost:5173/?video=<video_id>** ? rallies auto-load.
28 4. **Select ? Build ? download:** pick rallies ? Build ? **<video>** preview ? click the download link ? MP4 saved local.
29 5. **Repeat in Hindi:** click **???????** (or `?lang=hi`) ? redo select?compile?download ? confirm the whole flow works with the UI in Hindi.

30
31 ****Small unblocked SPA polish that makes Path A more "via the UI"** (independent of A/B, useful in both): a **"Load a video by ID"** input (paste the `video\_id` ? load, instead of editing the URL) + consume the engine's **download_url** (instead of hand-building it). One tight PR.

32
33 **## 4. Open decisions ? OWNER (these gate the dry run)**
34 1. **Scope:** Path A (cheapest dry run now) or Path B (the in-UI bucket ingest ? **SPA-only**; the engine endpoint already accepts `gs://`, see §1)? **Recommend A first** (proves the multilingual workflow fast), then B as the appliance build. ? ? **RESOLVED: A first (done, §7)**, then B (B-P2, the SPA ingest/progress screen).
35 2. **Compute / detector:** **Gemini** (real rallies; needs your **GEMINI_API_KEY**; uploads downsized chunks to Google = paid + cloud) . **motion** (key-free, runs anywhere, but lower-quality boundaries + fabricated server/shot metadata stays hidden) . **GPU WASB** (best quality, but needs a GPU host + is owner-parity-gated, not production-verified). **Recommend Gemini** for "real-ish, cheapest" ? pending your key + OK (paid + uploads footage).
36 3. **Engine host:** local Windows (simplest, Gemini/motion) . the CPU droplet `168.144.126.176` (shared testbed ? careful) . a GPU VM (only for WASB).
37 4. **Which bucket video:** e.g. `gs://khesutra-rally-corpus/videos/mahadevpura\_smoke\_180s.mp4` (already staged, known-good) or another clean ?720p clip (must pass the engine validators).

```

38
39     ## 5. Honest caveats
40     - **Gemini = paid + uploads footage to Google** (an external/paid action ? needs your
explicit OK + a clean clip). **Rotate the chat-shared `GEMINI_API_KEY`.
41     - **motion** fabricates `server`/`receiver`/`shot_count` (the SPA correctly hides them)
? boundaries are real but metadata isn't.
42     - **GPU/native WASB** is partially wired + owner-byte-parity-gated ? **not production-
verified**.
43     - Path A's "point at the bucket" is a **CLI pre-step**, not in-UI ? full in-UI ingest is
Path B.
44     - Edge auth: none in the engine (localhost dry-run is fine; required before any non-
localhost exposure).
45
46     ## 6. Definition of done (the dry run)
47     - [x] A bucket video ? (engine processes) ? the **SPA loads its rallies** ? select ?
compile ? **reel downloaded to the local machine**, all driven from the UI (Path A allows a one-
line CLI pre-process). ? ? **Path A done (local, 2026-06-21)**, see §7.
48     - [x] The **same flow works with the UI in Hindi** ? multilingual workflow confirmed. ?
? (Kannada also works via the same switcher.)
49     - [x] *(Path B = B-P2, SPA-only)* the bucket video is pointed-at **from the UI** (no CLI
pre-step). ? ? **DONE (live e2e, 2026-06-21), see §9.** In the SPA: entered a `gs://` URI ?
`POST /api/process` ? polled `/api/jobs` ? **12 rallies loaded** ? Build ? **reel downloaded
(1080p, 76s)**, all from the UI. Surfaced + fixed an engine compile bug along the way (engine
[#289](https://github.com/Khelsutra/badminton-highlight-indexer/pull/289)). **Screencast:**
`gs://khelsutra-rally-corpus/promo/khelsutra-bucket-cuj-screencast.mp4`.
50     - [ ] *(option 3)* the same CUJ confirmed on the **CPU droplet** (Linux deploy target) ?
droplet reachable but currently **un-provisioned** (no engine checkout; ~3.8GB RAM) ? needs a
from-scratch engine install or a proper deploy of the fixed engine.
51
52     ## 7. ? Run result ? Path A, local (2026-06-21)
53     Ran the real CUJ end to end on the local Windows box, isolated engine off
`origin/master` (fresh DB/output worktree), **Gemini** detector:
54     1. `gs://khelsutra-rally-corpus/videos/mahadevpura_smoke_180s.mp4` (1080p/30fps) ?
`gsutil cp` local ? `POST /api/process {segmenter: gemini}`.
55     2. Engine processed via Gemini (3 chunks) ? **12 rallies** (16 found, 4 too-short
rejected), `video_id 05e0e577?`.
56     3. SPA (`npm run dev`, Vite proxies `/api`??:8000) opened at `?video=05e0e577?` ?
**loaded the 12 real rallies** ("Loaded 12 rallies from this session").
57     4. **Build ? `/api/compile` ? reel** (1920x1080, 76.5s) ? **downloaded to
~/Downloads/khelsutra-dryrun/highlights_local_en.mp4** (50.9 MB).
58     5. **Repeated with the UI in Hindi** (`?lang=hi`): the page + rally cards + "??? ??????"
+ "???? ???? ????? ??" all in Hindi; reel compiled + **downloaded** (`highlights_local_hi.mp4`).
? multilingual workflow confirmed.
59
60     Screencast/proof: per-locale screenshots; both reels on the local machine. Screencast of
the i18n SPA is at `gs://khelsutra-rally-corpus/promo/khelsutra-i18n-screencast.mp4`.
61
62     ### ? Fix the dry run surfaced (engine) ? MERGED-PENDING
63     The async-job server **hung every job in `queued`** when run the documented way (`python
-m backend.main`): main.py runs as `__main__` (sets `__main__.db_instance`), but `job_worker`
reads `import backend.main` ? a **separate** module whose `db_instance` is `None` ? the drain
crashes silently. Latent because the unit tests monkeypatch the executor inline and the UI was
verified vs a stub engine. **Fix + a real-subprocess regression test ? engine
[#281](https://github.com/Khelsutra/badminton-highlight-indexer/pull/281)** (dispatch through
`backend.main`; FAILS-without / PASSES-with). Unblocked the run via the single-module
invocation; the fix applies to the droplet too.
64
65     ## 8. ? Correction (2026-06-21): B-P1 (the bucket-ingest endpoint) was already shipped ?
verified against code
66     **What this doc claimed (§1, original):** "`POST /api/process` accepts only a **local**
`video_path`" and "bridging `gcs://` ? a queryable `video_id` in the API DB is **net-new**
(T12/T14).".
67
68     **What the code actually does `[verified file:line, this session]`:** `POST
/api/process` has accepted bucket/Drive URIs since **engine M2

```

```

([#280](https://github.com/Khelsutra/badminton-highlight-indexer/pull/280))**. A
`gs://`/`gcs://`/`s3://`/`local://` `video_path`:
69 1. is resolved at submit (`backend/main.py:274` ? `storage.resolve_input_ref`;
`gs://`/`gcs` at `backend/storage/ref.py:60`; explicit scheme wins under default
`input_backend=local` at `ref.py:87`),
70 2. **skips the local-`os.path.exists` gate** for non-local backends (`main.py:283`) ?
returns **202, not 404**,
71 3. is fetched to scratch by the worker (`_execute_processing` ?
`storage.acquire_local_input`, `main.py:159` ? `GCSStorage.fetch_to_local`), then
72 4. runs the **same** pipeline as a local file, writing a queryable `video_id` to the API
SQLite DB.
73
74 **? Scope pin (default single-user mode):** this holds with
`security.allowed_video_root` **unset** (the local-appliance default). In **HOSTED mode**
(`allowed_video_root` set), a bucket URI is **correctly rejected 400** ? a tenant can't point
the engine at an arbitrary bucket. That is by design, not a regression.
75
76 **Already tested** (so this can't silently regress):
`tests/test_api_endpoints.py::test_process_nonlocal_uri_skips_exists_and_fetches`,
`?:test_process_hosted_mode_rejects_nonlocal_uri` (the hosted-mode counterweight),
`?:test_process_unknown_scheme_returns_400`,
`?:test_cli_nonlocal_input_fetches_and_keeps_uri`. So **B-P1 needs zero new engine feature
code** ? the only Path-B work is **B-P2 (the SPA ingest/progress screen)**. *(The hermetic test
fakes the network download via `m.storage.fetch`, so it proves the CODE routing, not the live
GCS fetch ? that needs `google-cloud-storage` + ADC on the box, verified at deploy time / option
3.)*
77
78 **Why it went stale (anti-spiral lesson):** the $1 claim was tagged `[verified]` but was
inferred from the **worker CLI** (`backend.worker.infer`, which writes CSV to the bucket) rather
than checked against the **API endpoint's** routing code ? and M2/#280 had since wired URI
ingest into `/api/process`. **The engine ships PRs between sessions, so its surface drifts.**
Convention going forward (also baked into the memory steering): **before scoping any
`[gap]`/`net-new` engine claim as work, re-verify it against engine code at a cited `file:line`
? a `[verified]` tag without a `file:line` is a hypothesis, not a fact.**
79
80 ## 9. ? Run result ? Path B, live in-UI bucket?reel CUJ (2026-06-21)
81 The full Path-B flow, driven **entirely from the UI** (no CLI pre-step), against a real
local engine (`python -m backend.main --server`, Gemini detector,
`GOOGLE_APPLICATION_CREDENTIALS` for the GCS fetch):
82 1. Opened the SPA ? the **"Process a video"** ingest panel (B-P2) ? pasted
`gs://khelsutra-rally-corpus/videos/mahadevpura_smoke_180s.mp4` ? **Process video**.
83 2. The engine **fetched the bucket video** (`google-cloud-storage` + ADC) ? **Gemini**
segmented it (live "Processing? segmenting (gemini)" progress in the UI) ? **12 rallies**
(`video_id 05e0e577?`), handed off via `?video=` ? the existing back-half rendered them.
84 3. **Build reel** ? `/api/compile` ? **reel preview + download** ? MP4 saved locally
(**1920x1080, 76s**).
85 4. **Screencast of the whole CUJ** (Playwright-recorded; the ~3-min Gemini wait fast-
forwarded): `gs://khelsutra-rally-corpus/promo/khelsutra-bucket-cuj-screencast.mp4`.
86
87 ### ? Bug the run surfaced + fixed (engine) ? MERGED
88 `/api/compile` **500'd** on a bucket-ingested video: `videos.filepath` keeps the
**source URI** (`gs://`, M2 provenance), and the stitcher was handed that directly ? ffmpeg
can't open a `gs://` path. **Fix:** compile now resolves the source via
`storage.acquire_local_input` first (identity for a local path; fetch-to-scratch for a bucket
URI) + a regression test ? engine **[#289](https://github.com/Khelsutra/badminton-highlight-
indexer/pull/289)** (full suite 1182?). This is what makes the in-UI bucket?reel CUJ actually
complete. *Proof-and-validation-driven: recording the screencast is what found the bug.*
89
90 ### Changelog
91 - `2026-06-21` ? created; Path-A run recorded ($7).
92 - `2026-06-21` ? **$8 correction:** verified B-P1 (`gs://`/`/api/process`?API-DB
`video_id`) is already shipped + tested (M2/#280); $1/$2/$4/$6 reframed so Path B = SPA-only
(B-P2). Sibling stale claims fixed in `LOCAL_APPLIANCE_ARCHITECTURE.md` (`/api/health` exists;
bucket-URI ingest not a net-new endpoint; T12 engine-half done) + `README.md`.
93 - `2026-06-21` ? **$9: Path-B live e2e DONE** ? the in-UI `gs://`?reel CUJ ran end to

```

end (12 rallies ? reel downloaded); surfaced + fixed the compile-from-bucket 500 (engine #289);
screencast uploaded to the promo bucket. **DoD §6 box 3 ?.**

94

5. assistant (2026-06-22T09:46:34.744Z)

My worklist is the two stale findings for `docs/DRY\_RUN\_CUJ.md`:

1. ****Status header stale**** (line 3): DRAFT/"Goal for this session" ? should be DONE/DELIVERED, since docs/README.md marks it DELIVERED and DoD §6 boxes 1-3 are checked.
2. ****en/kn/hi (3 langs) is stale**** (line 11): SPA now supports 6 (es/da/id added in #87).

Let me re-verify both against the live repos before editing.

6. user (2026-06-22T09:46:37.466Z)

50b2bd7 chore(repo): ignore local runtime artifacts + correct stale NEXT_STEPS claims (#94)
e2570e4 docs: archive the Friendly & Multilingual tracker (engineering complete) (#93)
0bef61e docs(friendly-multilingual): close out the chapter ? log #89/#90/#91 + strengthen fitNav tests (#92)
9f77ddc test(site): browser layout-regression CI guard + auto-collapse the nav to fit every locale (#91)
4415fe8 feat(site): localize remaining home sections (6 langs) + smooth the responsive header (#90)
---README---

khelsutra docs index

Tracked work (this repo)

- ****`DRY\_RUN\_CUJ.md`**** ? ? ****DELIVERED (2026-06-21):**** the UI drives the tool on a ****GCS-bucket video**** ? review/select ? compile ? ****reel downloaded local****, verified in en + Hindi. Path A (local) + Path B/B-P2 (in-UI `gs://` ingest, [#56]) both live e2e with a screencast (§9); the only remaining box (droplet, option 3) is ****parked by the owner**** (GPU subscription). Two engine bugs the runs surfaced are fixed (#281, #289). Was a milestone of `LOCAL\_APPLIANCE\_ARCHITECTURE.md`.

- ****`archives/FRIENDLY\_AND\_MULTILINGUAL.md`**** ? ? ****DONE ? ARCHIVED (2026-06-22).**** Made the product super user-friendly for non-tech-savvy coaches/parents ****and**** localized it to all ****6 languages**** across both surfaces (`web/` SPA ? react-i18next, `site/` ? a shared `t()` shim); the home-page body is fully localized (except the deferred `#join` form, A6) and the responsive header is fixed + ---i18n---

```
import daCommon from "../../locales/da/common.json";
import idCommon from "../../locales/id/common.json";

// Supported locales: en/kn/hi foundation (FRIENDLY D6) + MULTILINGUAL_EXPANSION Phase A
(es/da/id,
// Latin, catalog-only). Anything else falls back to `en`.
export const SUPPORTED_LOCALES = ["en", "kn", "hi", "es", "da", "id"] as const;
export type Locale = (typeof SUPPORTED_LOCALES)[number];
export const DEFAULT_LOCALE: Locale = "en";

// ONE persisted-preference key, shared with the marketing site's t() shim (D9 / tracker B3), so
a
// visitor's language choice carries across both surfaces.
```

7. user (2026-06-22T09:46:39.087Z)

cf59fbf feat(site): add es/da/id to the marketing site (parity with the SPA's 6 languages) (#87)
---grep locales---

```
25:export const SUPPORTED_LOCALES = ["en", "kn", "hi", "es", "da", "id"] as const;
26:export type Locale = (typeof SUPPORTED_LOCALES)[number];
```

52: supportedLngs: [...SUPPORTED_LOCALES],

8. assistant (2026-06-22T09:46:44.603Z)

Both findings are confirmed:

1. ****Status header****: `docs/README.md:4` marks this doc `? DELIVERED (2026-06-21)`; DoD §6 boxes 1-3 are `[x]`. The DRAFT/"Goal for this session" header is stale. The README also confirms the droplet (box 4) is owner-parked.
2. ****en/kn/hi****: `web/src/i18n/index.ts:25` confirms `SUPPORTED\_LOCALES = ["en", "kn", "hi", "es", "da", "id"]` (6 langs); #87 (cf59fbf) added es/da/id. Line 11's "Fully en/kn/hi" describing the **current** switcher is stale.

Now I'll apply the two surgical edits.

9. user (2026-06-22T09:46:47.487Z)

The file C:\Users\avidu\Projects\khelsutra\docs\DRY_RUN_CUJ.md has been updated successfully.
(file state is current in your context ? no need to Read it back)

10. user (2026-06-22T09:46:48.702Z)

The file C:\Users\avidu\Projects\khelsutra\docs\DRY_RUN_CUJ.md has been updated successfully.
(file state is current in your context ? no need to Read it back)

11. assistant (2026-06-22T09:46:55.294Z)

Both edits applied. Note the suggestedFix for the en/kn/hi finding explicitly said "the Hindi/Kannada run results in §7 stay accurate as historical record" ? I left §5/§7/§8/§9's historical "Hindi" run descriptions untouched, only updating the line-11 description of the **current** switcher capability, as scoped.

The third stale finding (line 15/68-72, engine line-number drift main.py:274/283/159 ? now 581/590/458) was rated by the audit as `correct` with severity `low` and an **optional** fix ("Optionally refresh the cited line numbers"). Since it's not in the staleFindings worklist (verdict was "correct", not "stale") and the routing logic itself is intact, I'm not touching it ? that would be scope creep beyond the assigned stale findings. The unverifiable GCS-artifact findings were correctly left alone.

Now returning the structured report.

12. user (2026-06-22T09:47:10.144Z)

Structured output provided successfully

13. assistant (2026-06-22T09:47:18.412Z)

I've completed the staleness fix for `docs/DRY\_RUN\_CUJ.md`. Both assigned stale findings were re-verified against the live repos and applied as surgical edits:

1. ****Status header (line 3)**** ? flipped from `DRAFT ? owner-gated . ... Goal for this session.` to `DONE ? DELIVERED (2026-06-21)` with the DoD §6 box-status noted and the droplet box flagged as owner-parked. Verified against `docs/README.md:4` (`? DELIVERED (2026-06-21)`) and the doc's own checked DoD boxes.
2. ****Language count (line 11)**** ? `Fully en/kn/hi` ? `Fully en/kn/hi/es/da/id (6 languages)`. Verified against `web/src/i18n/index.ts:25` (`SUPPORTED\_LOCALES = [en,kn,hi,es,da,id]`) and PR #87 (`cf59fbf`). I left the historical Hindi/Kannada run results in §5/§7/§9 untouched, per the audit's note.

Two items skipped (both correctly outside the stale worklist): an optional engine line-number refresh (audit verdict was `correct`, not stale ? refreshing brittle anchors would be scope creep) and the unverifiable GCS-artifact/run-result specifics (no bucket access). No code changes were needed.

The full structured report has been submitted.