

claude-windows / 8d4866bf-8edd-4da9-909e-022ddee12675

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\8d4866bf-8edd-4da9-909e-022ddee12675\subagents\workflows\wf_f9255361-69d\agent-ade92cb08be80423d.jsonl`  
- SHA-256: `c56c27d9b54de40d7d37f6c81c31aa168c0daae65fdb93792aa66922ef531bce`  
- Source modified: `2026-06-22T09:34:29+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `wf_f9255361-69d`  
- session_id: `8d4866bf-8edd-4da9-909e-022ddee12675`
```

Transcript

1. user (2026-06-22T09:32:46.035Z)

You are auditing ONE doc in the KhelSutra brand/site repo for staleness, as part of a "are all the docs up to date?" sweep.

Repo root: C:\Users\avidu\Projects\kheelsutra (bash path: /c/Users/avidu/Projects/kheelsutra)
The repo is on `master`, freshly synced to remote HEAD ? the latest merged PR is #94. Today is 2026-06-22.

DOC TO AUDIT: site/README.md

Your job: read the doc IN FULL, extract every concrete, falsifiable claim, and VERIFY each against the LIVE repo ? do NOT trust the doc's own wording. Hard discipline rule: a claim is "correct" ONLY when you have file:line or command-output evidence for it; if you cannot verify it, mark verdict "unverifiable" ? never guess, never assume the doc is right. Re-check against actual code/git.

Kinds of claims to check:

- File / path / route references ("served at /capture", "see src/store.js", "public/og.png", "systemd kheelsutra-site") -> does the path or route ACTUALLY exist now? Confirm with git ls-files / Read / the routing in site/server.js + site/src/index.js.
- PR numbers and counts ("PRs #1-#8 merged", "20 tests green", "6 languages", "3-step wizard", "P0..P5 done") -> verify against `git -C <root> log`, the actual test output/count, and the code.
- Status markers (DONE / pending / IN PROGRESS / DRAFT / "?" / "owner-gated" / "NEXT") -> still accurate vs the code & git history? Is something marked pending that is actually shipped, or vice-versa?
- Dates, "live at <url>", "merged in #N".
- Whether the doc describes something now SUPERSEDED or fully shipped -> is it an archive candidate (a tracker whose work is complete, or content duplicated by docs/archives/)? Note: docs/archives/FRIENDLY_AND_MULTILINGUAL.md already exists as the archived multilingual tracker.

Tools: Read the doc + every cited file; Grep/Glob to locate code; Bash for `git -C /c/Users/avidu/Projects/kheelsutra log --oneline`, `... ls-files`, and greps. To check a test count, you may run the suite (`cd /c/Users/avidu/Projects/kheelsutra/site && node --test` or `npm test`), but prefer cheap checks. Verify against master as it stands NOW.

Severity: high = a factually wrong claim a reader would act on (wrong path/route, a "DONE" that isn't, a wrong count/PR that misleads). medium = outdated-but-harmless (stale PR range, old "where we are", superseded framing). low = cosmetic.

Do NOT modify any file ? audit only. Return ONLY the structured object (doc, overallStatus, findings[], summary).

2. assistant (2026-06-22T09:32:49.780Z)

I'll audit `site/README.md` against the live repo. Let me start by reading the doc and gathering repo context.

3. user (2026-06-22T09:32:50.730Z)

```
1      # site/ ? public marketing site + Request-Access API (Cloudflare Worker)
2
3      The public site at **[kheelsutra.guru](https://kheelsutra.guru)** ? a static landing page
+
4      `/capture` checklist ? plus a small Worker that backs the Request-Access form with
**Cloudflare D1**.
5
6      ## Layout
7      - `public/` ? static assets served as-is (`index.html`, `capture.html`, `og.png`,
`robots.txt`).
8      - `src/index.js` ? the Worker: serves `public/` and handles `POST /api/request-access` ?
D1.
9      - `src/index.test.js` ? unit tests for the form logic.
10     - `wrangler.toml` ? Worker + static-assets + D1 binding config.
11     - `schema.sql` ? the `access_requests` D1 table.
12     - `make_og.py` ? regenerates `public/og.png` (Pillow + Windows fonts): `python
make_og.py`.
13
14     ## Develop / test
15     ```
16     npm test      # node --test src/ (no install needed)
17     npm run dev   # npx wrangler dev ? local Worker + D1
18     ```
19
20     ## Deploy (owner ? our Cloudflare account)
21     One-time:
22     ```
23     npx wrangler d1 create khelsutra-access          # paste database_id
into wrangler.toml
24     npx wrangler d1 execute khelsutra-access --remote --file=schema.sql
25     ```
26     Then either:
27     - **Manual:** `npx wrangler deploy`, or
28     - **Git (recommended):** connect this repo in Cloudflare ? Workers Builds (root dir
`site/`), bind the
29     D1 database ? every merge to `master` deploys; PRs get preview URLs.
30
31     ## Notes
32     - The form posts to `/api/request-access`; if the Worker is unreachable it falls back to
a prefilled
33     `mailto:hello@`, so it degrades gracefully.
34     - `hello@kheelsutra.guru` relies on Cloudflare Email Routing.
35     - The Web Analytics beacon in `public/index.html` has a placeholder token
(`REPLACE_WITH_CF_TOKEN`)
36     ? paste the real token or enable Cloudflare automatic setup.
37
38     ## Crawling / bot protection
39
40     `kheelsutra.guru` sits behind Cloudflare's edge security (Bot Fight Mode / WAF / rate
limiting), which runs **in front of** the Worker. The app itself never returns `403` ?
41     `src/index.js` only emits `405`/`400` on the form path and `200`/`404` for assets ? so
42     **a `403` while crawling is Cloudflare's edge, not this site**. `public/robots.txt` is
43     advisory (and won't change a `403`); the edge mitigation is the thing to adjust.
44
45
46     ### Diagnose which Cloudflare feature is blocking
47     From the crawler's host (prints status + Cloudflare headers):
48     ```
49     curl -sS -D - -o /dev/null -A "<crawler UA>" "https://kheelsutra.guru/<failing path>"
50     ```
51     - `cf-mitigated: challenge` (or a "Checking your browser" body) ? **Bot Fight Mode**
```

```

52     - body "Sorry, you have been blocked" + a Ray ID ? **WAF rule**
53     - HTTP `429`, or `403` + a `Retry-After` header ? **Rate limiting**
54
55     Authoritative: **Cloudflare dashboard ? khelsutra.guru ? Security ? Events**, filter
56     Action = Block/Challenge ? each row names the blocking service + the Ray ID.
57
58     ### Let your OWN crawler through (allowlist; don't weaken protection globally)
59     1. **IP Access Rule = Allow** (Security ? WAF ? Tools ? IP Access Rules) for the
crawler's
60         egress IP (`curl -s https://api.ipify.org` from the crawler). Bypasses Bot Fight Mode
+
61         security level; **works on the free plan** ? the simplest reliable fix.
62     2. **WAF Skip rule by secret header** (Pro+): custom rule
63         `any(http.request.headers["x-crawl-key"][*] eq "<secret>")` ? Action **Skip ? Super
Bot
64         Fight Mode + managed rules + rate limiting**; have the crawler send `X-Crawl-Key:
<secret>`.
65         Robust regardless of IP/UA. (Plain Bot Fight Mode on the free plan can't be
selectively
66         skipped ? use #1 or toggle it off.)
67     3. **Toggle Bot Fight Mode off** (Security ? Bots) for a short crawl window.
68     4. **Rate-limited?** Lower concurrency + add a delay (~1 req/s); a burst of identical-
looking
69         requests trips heuristics even with a good UA.
70
71     Crawler-side (no Cloudflare change): send a real browser `User-Agent` + `Accept` /
72     `Accept-Language`. Light crawling may pass, but Bot Fight Mode fingerprints beyond the
UA,
73     so #1/#2 are the durable fixes.
74

```

4. user (2026-06-22T09:32:51.411Z)

```

50b2bd7 chore(repo): ignore local runtime artifacts + correct stale NEXT_STEPS claims (#94)
e2570e4 docs: archive the Friendly & Multilingual tracker (engineering complete) (#93)
0bef61e docs(friendly-multilingual): close out the chapter ? log #89/#90/#91 + strengthen fitNav
tests (#92)
9f77ddc test(site): browser layout-regression CI guard + auto-collapse the nav to fit every
locale (#91)
4415fe8 feat(site): localize remaining home sections (6 langs) + smooth the responsive header
(#90)
26c9da1 feat(site): localize the hero demo mockup + #how feature block (6 languages) (#89)
7ae9cee feat(web): P2c first-run setup wizard (3 steps from /api/capabilities, 6 languages)
(#88)
cf59fbf feat(site): add es/da/id to the marketing site (parity with the SPA's 6 languages) (#87)
17242a9 feat(site): localize own-model body content to kn/hi (#86)
7b04f02 deploy(kit): align with the live-verified droplet deploy (+ the gemini-default fix)
(#85)
5e5a2ac docs(friendly-multilingual): close out the chapter ($9 DoD + UI/UX + B7 deferred) (#84)
4c0f35c fix(site): UI/UX ? compact language switcher + header overflow fix + localized hero
(#83)
0fae364 feat(web+site): A2 ? localized WhatsApp share for the reel (SPA + site) (#82)
1e52cfd deploy(droplet): tailored alpha kit for the CPU droplet (engine api. side) (#81)
64076a5 test(web): add a coverage floor (CI-gated) + i18n regression tests for the #77 audit
(#80)
8328721 docs(tracker): refresh the $7 milestone prose to merged reality (A6/A7/#77 landed) (#79)
3f0cdb8 docs(deploy): fix two would-break-the-deploy gaps in the alpha runbook
(RALLY_ALLOWED_HOSTS + gs:// jail) (#78)
9f1ec9c fix(i18n): native-review pass on the kn/hi SPA console catalogs (#77)
9a2d7d2 feat(site): A7 ? camera-placement diagram on the capture page (#76)
aa8d5e9 feat(site): A6 ? plain-language Request-Access GPU/storage questions (#75)

```

5. user (2026-06-22T09:32:52.118Z)

site/README.md

```
site/make_og.py
site/package-lock.json
site/package.json
site/public/capture.html
site/public/flyer-hi.html
site/public/flyer-kn.html
site/public/flyer.html
site/public/i18n.js
site/public/img/annotate-enable.png
site/public/img/annotate-tool.png
site/public/img/flyer-qr.svg
site/public/img/studio-en.png
site/public/img/studio-hi.png
site/public/img/studio-kn.png
site/public/index.html
site/public/locales/da/common.json
site/public/locales/en/common.json
site/public/locales/es/common.json
site/public/locales/hi/common.json
site/public/locales/id/common.json
site/public/locales/kn/common.json
site/public/og.png
site/public/own-model.html
site/public/robots.txt
site/public/video/tool-cuj.mp4
site/schema.sql
site/scripts/copy-locales.mjs
site/scripts/droplet-smoke.sh
site/scripts/parity-check.sh
site/scripts/provision-vm.sh
site/scripts/teardown-vm.sh
site/server.js
site/src/assets.test.js
site/src/deploy-config.test.js
site/src/i18n.shim.test.js
site/src/index.js
site/src/index.test.js
site/src/layout.browser.mjs
site/src/locales-drift.test.js
site/src/notify.js
site/src/server.test.js
site/src/store.js
site/src/ux.dom.test.js
site/wrangler.toml
```

6. assistant (2026-06-22T09:32:55.384Z)

The doc references several files and routes. Let me verify the key claims: the file layout, the Worker route paths, the test command, and check whether `server.js`/`store.js`/`notify.js` (which exist but aren't mentioned) change the picture. Let me read the cited files.

7. user (2026-06-22T09:32:56.343Z)

```
1      // Cloudflare Worker for khelsutra.guru ? serves the static site and handles the
2      // Request-Access form: POST /api/request-access -> validate -> insert into D1 (binding
DB).
3      // Pure helpers are exported so they can be unit-tested without a live Worker/D1.
4
5      import { d1Store } from "./store.js";
6      import { cfEmailNotify } from "./notify.js";
7
8      const SEGMENTS = ["academy", "individual"];
9      const GPU = ["yes", "unsure", "no", "prefer-cloud"];
10     const COURTS = ["1", "2-4", "5+"];
11     const STORAGE = ["gdrive", "onedrive", "dropbox", "bucket", "local", "unsure"];
```

```

12
13 function str(v, max) {
14     if (v == null) return null;
15     const s = String(v).trim();
16     return s ? s.slice(0, max) : null;
17 }
18
19 // Honeypot: humans never fill the hidden "company" field.
20 export function isBot(p) {
21     return !! (p && typeof p.company === "string" && p.company.trim() !== "");
22 }
23
24 export function validate(p) {
25     if (!p || typeof p !== "object") return { ok: false, error: "bad payload" };
26     const email = String(p.email || "").trim();
27     if (!/^[^@\s]+@^[@\s]+\.[^@\s]+$/ .test(email)) return { ok: false, error: "valid email
required" };
28     return { ok: true };
29 }
30
31 export function toRow(p, meta = {}) {
32     const storage = Array.isArray(p.storage)
33         ? p.storage.filter((s) => STORAGE.includes(s)).join(",")
34         : "";
35     return {
36         segment: SEGMENTS.includes(p.segment) ? p.segment : null,
37         name: str(p.name, 200),
38         email: str(p.email, 320),
39         city: str(p.city, 200),
40         org_name: str(p.org_name, 200),
41         courts: COURTS.includes(p.courts) ? p.courts : null,
42         has_gpu: GPU.includes(p.has_gpu) ? p.has_gpu : null,
43         storage,
44         message: str(p.message, 4000),
45         consent: p.consent ? 1 : 0,
46         country: str(meta.country, 8),
47         user_agent: str(meta.userAgent, 400),
48     };
49 }
50
51 function json(obj, status = 200) {
52     return new Response(JSON.stringify(obj), { status, headers: { "content-type":
"application/json" } });
53 }
54
55 export async function handleRequest(request, store, notify) {
56     let body;
57     try {
58         body = await request.json();
59     } catch {
60         return json({ ok: false, error: "invalid json" }, 400);
61     }
62     if (isBot(body)) return json({ ok: true }); // silently drop bots
63     const v = validate(body);
64     if (!v.ok) return json({ ok: false, error: v.error }, 400);
65
66     const r = toRow(body, {
67         country: request.headers.get("cf-ipcountry"),
68         userAgent: request.headers.get("user-agent"),
69     });
70
71     await store.insertAccessRequest(r);
72     try { await notify?.notify(r); } catch { /* email is best-effort; never fail the
request */ }
73

```

```

74     return json({ ok: true });
75 }
76
77 export default {
78   async fetch(request, env) {
79     const url = new URL(request.url);
80     if (url.pathname === "/api/request-access") {
81       if (request.method !== "POST") return json({ ok: false, error: "method not
allowed" }, 405);
82       return handleRequest(request, d1Store(env.DB), cfEmailNotify(env.NOTIFY,
env.NOTIFY_FROM, env.NOTIFY_TO));
83     }
84     return env.ASSETS.fetch(request); // everything else: static assets
85   },
86 };
87

```

8. user (2026-06-22T09:32:56.826Z)

```

1  {
2    "name": "khelsutra-site",
3    "private": true,
4    "type": "module",
5    "scripts": {
6      "test": "node --test",
7      "test:cf": "node --test src/index.test.js",
8      "test:vm": "node --test src/server.test.js",
9      "test:deploy": "node --test src/deploy-config.test.js src/assets.test.js",
10     "test:layout": "node --test src/layout.browser.mjs",
11     "copy:locales": "node scripts/copy-locales.mjs",
12     "predeploy": "node scripts/copy-locales.mjs",
13     "deploy": "wrangler deploy",
14     "dev": "wrangler dev",
15     "og": "python make_og.py"
16   },
17   "devDependencies": {
18     "jsdom": "^29.1.1",
19     "playwright": "^1.61.0"
20   }
21 }
22

```

9. user (2026-06-22T09:32:57.149Z)

```

1  # Cloudflare Worker for khelsutra.guru ? static site + Request-Access form API.
2  #
3  # One-time setup (owner, our Cloudflare account):
4  #   npx wrangler d1 create khelsutra-access          # copy the printed database_id
below
5  #   npx wrangler d1 execute khelsutra-access --remote --file=schema.sql
6  #   npx wrangler deploy                             # or connect this repo via Workers
Builds (git)
7  #
8  name = "hidden-king-khelsutra" # the git-connected Worker that owns khelsutra.guru ?
keep in sync with Cloudflare
9  main = "src/index.js"
10 compatibility_date = "2025-06-01"
11
12 [assets]
13 directory = "./public"
14 binding = "ASSETS"
15
16 [[d1_databases]]
17 binding = "DB"
18 database_name = "khelsutra-access"

```

```

19     database_id = "cd02c6bf-4bd6-4a33-a333-ced86e36b511"    # from `wrangler d1 create
khelsutra-access`
20
21     # Email-on-registration: Email Routing send binding -> a VERIFIED destination address.
22     [[send_email]]
23     name = "NOTIFY"
24     destination_address = "avi.dullu@gmail.com"
25
26     [vars]
27     NOTIFY_FROM = "noreply@khelsutra.guru"
28     NOTIFY_TO = "avi.dullu@gmail.com"
29

```

10. assistant (2026-06-22T09:33:01.607Z)

The README describes the Worker. But there's clearly a second deployment surface (`server.js`, VM scripts, `notify.js`). Let me verify the specific README claims: public assets exist, `/capture` route, the form fallback to mailto, the analytics token, and check what `server.js` is (it may be a parallel VM server not mentioned). Let me also run the test count.

11. user (2026-06-22T09:33:03.315Z)

```

site/public/capture.html
site/public/flyer-hi.html
site/public/flyer-kn.html
site/public/flyer.html
site/public/i18n.js
site/public/img/annotate-enable.png
site/public/img/annotate-tool.png
site/public/img/flyer-qr.svg
site/public/img/studio-en.png
site/public/img/studio-hi.png
site/public/img/studio-kn.png
site/public/index.html
site/public/locales/da/common.json
site/public/locales/en/common.json
site/public/locales/es/common.json
site/public/locales/hi/common.json
site/public/locales/id/common.json
site/public/locales/kn/common.json
site/public/og.png
site/public/own-model.html
site/public/robots.txt
site/public/video/tool-cuj.mp4
---ROBOTS---

```

robots.txt ? khelsutra.guru (Cloudflare Worker: static site + Request-Access form API)

#

Crawling the marketing pages is welcome. The only non-asset path is the form API (`/api/request-access`, POST-only ? returns 405 to GET), which is not a crawl target.

#

NOTE: edge bot protection (Cloudflare Bot Fight Mode / WAF / rate limiting) runs in FRONT of this file and can still 403 automated clients regardless of these rules ? robots.txt is advisory, not enforcement. To crawl your OWN site through that layer, allowlist the crawler in Cloudflare (see site/README.md ? "Crawling / bot protection")

```

User-agent: *
Allow: /
Disallow: /api/

```

Advisory pacing for crawlers that honor it (Googlebot ignores Crawl-delay and paces itself). Keep automated traffic gentle ? this is a single small Worker.

Crawl-delay: 1

12. user (2026-06-22T09:33:05.088Z)

```
52:  /* in-flight spinner for the request-access submit */
393:      <a class="btn btn-dark" href="/capture" data-i18n="site.record.checklistCta">Full
1-page capture checklist ?</a>
620:      <div><a href="/own-model">Own your model</a> · <a href="/capture">Capture
checklist</a> · <a href="mailto:hello@khelsutra.guru">hello@khelsutra.guru</a> · <a
href="https://khelsutra.guru">khelsutra.guru</a></div>
629:<script defer src="https://static.cloudflareinsights.com/beacon.min.js" data-cf-
beacon='{ "token": "REPLACE_WITH_CF_TOKEN" }'></script>
717:      return 'mailto:hello@khelsutra.guru?subject=' + encodeURIComponent('KhelSutra alpha
access') +
738:      var r = await fetch('/api/request-access', {method:'POST', headers: {'content-
type': 'application/json'}, body: JSON.stringify(p)});
```

13. user (2026-06-22T09:33:08.034Z)

```
? [VM host] serves the tool screencast as video/mp4 (plays in the post-signup reveal) (3.7097ms)
? [VM host] serves /own-model (8.0552ms)
? [VM host] serves the academy flyers (/flyer, /flyer-kn, /flyer-hi) as clean URLs (10.3855ms)
? [VM host] POST /api/request-access persists to SQLite (17.2155ms)
? [VM host] honeypot + bad email are rejected without persisting (5.1625ms)
? [VM host] GET /api/request-access -> 405 (1.2569ms)
? [VM host · N8] a Kannada browser (Accept-Language: kn) gets the kn-rendered page (2.664ms)
? [VM host · N8] ?lang= overrides the browser language (2.5512ms)
? [VM host · N8] the persisted cookie beats Accept-Language (returning user keeps their choice)
(2.4152ms)
? [VM host · N8] unsupported/absent language falls back to en (never a locale we can't serve)
(4.0475ms)
? [VM host · N8] non-localized assets are NOT rewritten (only the 3 marketing pages negotiate)
(1.8928ms)
? [VM host] serves the shared catalogs the shim fetches (/locales/*/common.json as JSON)
(2.4874ms)
? submit shows an obvious in-flight state (disabled + aria-busy + 'Sending?') (359.0789ms)
? on failure the button is restored ? never stuck on a frozen 'Sending?' (80.676ms)
? after sign-up BOTH videos appear and a ? closes the reveal (69.0024ms)
? A2: after sign-up the reveal offers a WhatsApp share of the reel (wa.me, safe, localizable)
(42.4567ms)
? the auto-close HOLDS while the reel is playing (no mid-watch cut-off) (43.3387ms)
? tests 86
? suites 0
? pass 86
? fail 0
? cancelled 0
? skipped 0
? todo 0
? duration_ms 2232.9717
```

14. assistant (2026-06-22T09:33:15.090Z)

The README has verified accurately so far, but there's a significant gap: the test suite includes a whole `[VM host]` test surface (`server.js` ? an Express/Node VM server serving the SAME site, plus the SPA tool) that the README does not mention at all. The README presents the site as *only* a Cloudflare Worker. Let me confirm `server.js` is a real parallel deployment surface and check what `make\_og.py`/`schema.sql` claims hold.

15. user (2026-06-22T09:33:16.053Z)

```
1      // Always-on VM host for the KhelSutra site ? dependency-free (Node 22+ only).
```

```

2      // Serves public/ and POST /api/request-access into SQLite (node:sqlite), using the
3      // SAME handler logic as the Cloudflare Worker (src/index.js). This is the portability
4      // fallback "if Cloudflare acts out" ? see docs/PORTABILITY.md.
5      //
6      //   PORT=8080 DB_PATH=/var/lib/khelsutra/access.db node server.js
7      //
8      import { createServer } from "node:http";
9      import { readFile } from "node:fs/promises";
10     import { join, extname, normalize } from "node:path";
11     import { DatabaseSync } from "node:sqlite";
12     import { handleRequest } from "../src/index.js";
13     import { sqliteStore } from "../src/store.js";
14     import { noopNotify } from "../src/notify.js";
15
16     const PORT = Number(process.env.PORT) || 8080;
17     const ROOT = import.meta.dirname;
18     const PUBLIC = join(ROOT, "public");
19     const DB_PATH = process.env.DB_PATH || join(ROOT, "access.db");
20
21     const db = new DatabaseSync(DB_PATH);
22     db.exec(await readFile(join(ROOT, "schema.sql"), "utf8")); // CREATE TABLE IF NOT EXISTS
? idempotent
23     const store = sqliteStore(db);
24     const notify = noopNotify(); // VM host: no email yet (add SMTP later ? see
docs/PORTABILITY.md)
25
26     const TYPES = {
27       ".html": "text/html; charset=utf-8", ".png": "image/png", ".svg": "image/svg+xml",
28       ".css": "text/css", ".js": "text/javascript", ".json": "application/json;
charset=utf-8",
29       ".ico": "image/x-icon", ".txt": "text/plain",
30       ".mp4": "video/mp4", ".webm": "video/webm", ".gif": "image/gif", ".jpg": "image/jpeg",
".jpeg": "image/jpeg",
31     };
32
33     // --- Language negotiation (tracker B3 / $6.5 N8)
-----
34     // The trilingual marketing site. The client shim (public/i18n.js) does the full
negotiation in the
35     // browser; this VM host PRE-RESOLVES the locale server-side so the first paint (and any
no-JS read)
36     // already carries the right <html lang>. Same ORDER as the SPA + shim: persisted pref
-> ?lang= ->
37     // navigator/Accept-Language -> en. (On Cloudflare, assets are served directly and only
the client
38     // shim negotiates ? see src/index.js.)
39     const SITE_LOCALES = ["en", "kn", "hi", "es", "da", "id"];
40     const LANG_COOKIE = "khelsutra.lang";

```

16. user (2026-06-22T09:33:19.472Z)

```

-- D1 schema for Request-Access submissions.
-- Apply: npx wrangler d1 execute khelsutra-access --remote --file=schema.sql
CREATE TABLE IF NOT EXISTS access_requests (
  id          INTEGER PRIMARY KEY AUTOINCREMENT,
  created_at  TEXT NOT NULL DEFAULT (datetime('now')),
  ---MAKEOG---
1:"""Generate og.png (1200x630) ? the social-share card for khelsutra.guru.
5:from PIL import Image, ImageDraw, ImageFont
21:def font(name, size):
24:bold = lambda s: font("segoeui.ttf", s)
25:reg = lambda s: font("segoeui.ttf", s)
31:d.text((M, 90), "Khe", font=wf, fill=WHITE)
32:khel_w = d.textlength("Khe", font=wf)
36:d.text((tri_x + 44, 90), "Sutra", font=wf, fill=LIME)

```

```
39:d.text((M, 232), "Every rally, indexed.", font=bold(78), fill=WHITE)
40:d.text((M, 322), "The dead time, gone.", font=bold(78), fill=LIME)
---SRCFILES---
site/src/assets.test.js
site/src/deploy-config.test.js
site/src/i18n.shim.test.js
site/src/index.js
site/src/index.test.js
site/src/layout.browser.mjs
site/src/locales-drift.test.js
site/src/notify.js
site/src/server.test.js
site/src/store.js
site/src/ux.dom.test.js
```

17. assistant (2026-06-22T09:33:25.377Z)

All the README's explicit claims check out individually. The main staleness is one of ****omission/framing****: the README presents ``src/index.js`` as "the Worker" that "serves ``public/``" ? but per ``index.js:84``, on Cloudflare the Worker does NOT serve ``public/``; ``env.ASSETS.fetch`` does (Cloudflare static assets). And there's an entire parallel VM-host deployment surface (``server.js``, the live droplet at ``app.khelsutra.guru``, the SPA tool) that the README is silent on. Let me verify the README's two specific descriptive claims about `index.js` precisely, and check the "Layout" file list against ``src/`` reality.

18. user (2026-06-22T09:33:27.219Z)

```
f7e2bd5 chore(site): add robots.txt + a "crawling / bot protection" runbook for khelsutra.guru (#71)
f699847 feat(site): Request-Access Worker + D1 backend (+ 10 unit tests) (#3)
5d6eb83 feat(site): add the khelsutra.guru marketing landing page (#1)
---LAST MOD---
f7e2bd5 2026-06-21 23:20:54 +0530 chore(site): add robots.txt + a "crawling / bot protection"
runbook for khelsutra.guru (#71)
```

19. user (2026-06-22T09:33:29.328Z)

```
docs/PORTABILITY.md
PORTABILITY exists
README: no mention of server.js/VM host/droplet/PORTABILITY
```

20. assistant (2026-06-22T09:33:36.038Z)

This confirms the staleness picture. The README was last meaningfully updated at #71 (the "Develop/test" + layout sections date back to #3). Since then (#75-#94), the site gained: a whole VM-host deployment (``server.js``), multilingual support (6 languages, locales/), the SPA tool screencast, flyers, own-model page, and the ``src/`` directory grew from 2 files to 11. The README's "Layout" section and "npm test" description are now incomplete/imprecise.

Let me verify the two remaining precise claims: (1) ``npm test`` # `node --test src/`` ? the actual script is ``node --test`` (no ``src/`` arg), and (2) the layout list omits many ``public/`` assets and ``src/`` files.

21. user (2026-06-22T09:33:38.085Z)

```
README says: 'npm test # node --test src/'
package.json test script:
  "test": "node --test",
---
README claims index.js 'serves public/'. index.js line 84:
  return env.ASSETS.fetch(request); // everything else: static assets
```

22. assistant (2026-06-22T09:33:48.090Z)

I have everything I need. Let me compile the findings.

Summary of verification:

- **Layout list** (lines 7-11): Materially incomplete. ``public/`` now also has ``flyer*.html``, ``i18n.js``, ``img/``, ``locales/`` (6 langs), ``video/``, ``own-model.html`` ? none listed. ``src/`` now has 11 files (``server.js``, ``store.js``, ``notify.js``, plus 7 test files), but the README lists only ``index.js``/``index.test.js`` and omits ``server.js``, ``store.js``, ``notify.js`` entirely. Stale.
- **"npm test # node --test src"** (line 16): The actual script is ``node --test`` (not ``node --test src/``). Minor wrong-but-harmless.
- **"src/index.js" serves "public/"** (line 8): On Cloudflare the Worker delegates to ``env.ASSETS.fetch`` (line 84), not the Worker code itself. Imprecise but defensible framing.
- **Whole VM-host surface unmentioned**: ``server.js`` is a real second deployment (live on the droplet at ``app.khelsutra.guru`` per the test suite + memory) using the same handler. README frames the site as Cloudflare-Worker-only. Significant omission.
- **Verified-correct claims**: ``/capture`` route + ``capture.html`` exist; mailto fallback to ``hello@khelsutra.guru`` (line 717); ``REPLACE_WITH_CF_TOKEN`` analytics placeholder (line 629); ``robots.txt`` advisory; ``schema.sql`` ? ``access_requests`` table; ``make_og.py`` regenerates ``og.png`` via Pillow; ``wrangler.toml`` D1 binding; bot-protection runbook accurate; form posts to ``/api/request-access`` POST?405 on GET.

23. user (2026-06-22T09:34:18.221Z)

Structured output provided successfully

24. assistant (2026-06-22T09:34:29.346Z)

Audit complete. I verified ``site/README.md`` against the live repo at master HEAD (#94).

Overall status: **minor-staleness** (not an archive candidate ? it's a live operational README, not a completed tracker).

Key findings:

- **All explicit falsifiable claims verify correct**: the ``/capture`` route, mailto fallback (``site/public/index.html:717,738``), ``REPLACE_WITH_CF_TOKEN`` placeholder (``:629``), ``access_requests`` D1 table, ``make_og.py`` Pillow regen, and the entire bot-protection runbook (Worker only emits 405/400/200/404, never 403 ? ``site/src/index.js:60,64,81,84``).
- **Staleness is by omission/imprecision** ? the layout/dev sections date to #1/#3/#71 and haven't kept up through #94:
 - The whole **VM-host deployment surface is unmentioned**: ``site/server.js`` (an always-on Node host sharing the same handler, the live alpha-droplet path, ``docs/PORTABILITY.md``), plus ``src/store.js``/``src/notify.js``. README frames the site as Cloudflare-Worker-only.
 - **Layout under-describes "src/"** (lists 2 of 11 files) and **"public/"** (omits 6-language ``locales/``, ``flyers``, ``own-model.html``, ``img/``, ``video/``).
 - **"npm test" annotation is wrong**: README says ``node --test src/``; actual script is ``node --test`` (runs 86 tests).
 - **"src/index.js serves "public/"** is imprecise ? on Cloudflare the ``ASSETS`` binding serves statics; the Worker delegates (``site/src/index.js:84``).

Full structured findings (10 entries with file:line evidence and suggested fixes) returned via StructuredOutput.