

claude-windows / 188dd06f-9c99-487c-aa7d-b0bee83c629c

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\188dd06f-9c99-487c-aa7d-b0bee83c629c\subagents\workflows\wf_416cc0e5-a35\agent-aae52d12d479e581f.jsonl`
- SHA-256: `68110fcfa6a445683527fb37b9167fd7353fd59764b85f9ab7a64f788d9897c9`
- Source modified: `2026-06-22T09:50:45+00:00`
- Imported at: `2026-07-05T16:48:28+00:00`
- project: `wf_416cc0e5-a35`
- session_id: `188dd06f-9c99-487c-aa7d-b0bee83c629c`
```

Transcript

1. user (2026-06-22T09:47:32.509Z)

Review the NIGHTLY REEL-COUNT engine-regression change under C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/nightly-e2e (a NEW, additive PR). It runs the configured WASB pipeline on a GCP GPU VM nightly, asserts reel-count (==) + quality metrics + cost, emails a report, self-deletes. Files: scripts/nightly_reelcount.py, scripts/nightly_email.py, deploy/gcp/nightly_regression/run_on_vm.sh, deploy/gcp/nightly_regression/launch.sh, deploy/gcp/nightly_regression/register_scheduler.sh, deploy/gcp/nightly_regression/Dockerfile.launcher, eval_baselines/reelcount_manifest.json. Read them under that absolute path. Cite file:line. Report only REAL defects (correctness/safety/security that would bite in production or a regression that escapes), not style. Return raw structured findings.

ADVERSARIALLY VERIFY this finding ? confirm with concrete code evidence under C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/nightly-e2e or REFUTE it (default real=false if the evidence isn't clearly there). Finding:

```
{
  "title": "SMTP password passed to smtplib.SMTP.login() without connection security verification",
  "file": "scripts/nightly_email.py",
  "line": "86-89",
  "severity": "low",
  "detail": "The code uses `s.starttls()` (line 87) before login, which is correct. However, there is no explicit certificate verification context passed to SMTP(). By default, Python's smtplib uses the system CA store, but the code does not explicitly set `context=ssl.create_default_context()` to enforce VERIFIED mode. An MITM attacker could potentially intercept and downgrade the connection (though STARTTLS mandatory mode should prevent this). This is a minor issue because Gmail's STARTTLS is strict.",
  "fix": "Explicitly pass `context=ssl.create_default_context()` to smtplib.SMTP() to enforce certificate verification: `with smtplib.SMTP(host, port, timeout=60, context=ssl.create_default_context()) as s:`"
}
```

2. assistant (2026-06-22T09:47:34.616Z)

I'll analyze the NIGHTLY REEL-COUNT engine-regression change by examining the relevant files, with a focus on verifying the SMTP security finding.

3. user (2026-06-22T09:47:35.925Z)

```
1      """Email the nightly regression report ? stdlib SMTP, creds from env or GCP Secret
Manager.
2
3      Used by ``scripts/nightly_reelcount.py --email`` (and reusable by any nightly).
Deliberately
4      dependency-light (stdlib ``smtplib`` + ``email.message``) and **graceful**: if no SMTP
```

```

password can
5         be resolved it logs + returns ``False`` (the report is still written/printed ? a broken
mailer must
6         never mask a regression).
7
8         Credential resolution (first hit wins), so it works on a GCP VM, locally, or in tests:
9         ? Recipient: ``--email-to`` arg ? ``NIGHTLY_EMAIL_TO`` ? ``avi.dullu@gmail.com``
10        ? SMTP host: ``NIGHTLY_SMTP_HOST`` (default ``smtp.gmail.com``) . port
``NIGHTLY_SMTP_PORT`` (587, STARTTLS)
11        ? SMTP user: ``NIGHTLY_SMTP_USER`` (default = recipient)
12        ? SMTP pass: ``NIGHTLY_SMTP_PASS`` ? Secret Manager ``NIGHTLY_SMTP_SECRET``
13                (`projects/<p>/secrets/<name>/versions/latest``; via the google-cloud-
secret-manager
14                lib if installed, else the ``gcloud`` CLI). A Gmail **app password** is
the simplest.
15
16        Set the secret once (owner): echo -n '<app-password>' | gcloud secrets create nightly-
smtp-pass --data-file=-
17        and on the VM: export NIGHTLY_SMTP_SECRET=projects/<proj>/secrets/nightly-smtp-
pass/versions/latest
18                export NIGHTLY_SMTP_USER=<your-gmail>
NIGHTLY_EMAIL_TO=avi.dullu@gmail.com
19        ""
20        from __future__ import annotations
21
22        import logging
23        import os
24        from email.message import EmailMessage
25        from typing import Optional
26
27        log = logging.getLogger("nightly_email")
28
29        DEFAULT_TO = "avi.dullu@gmail.com"
30        DEFAULT_HOST = "smtp.gmail.com"
31        DEFAULT_PORT = 587
32
33
34        def build_message(subject: str, body_md: str, from_addr: str, to_addr: str) ->
EmailMessage:
35        """A text/plain email carrying the markdown report (renders fine in any client).
Pure ? testable."""
36        msg = EmailMessage()
37        msg["Subject"] = subject
38        msg["From"] = from_addr
39        msg["To"] = to_addr
40        msg.set_content(body_md)
41        return msg
42
43
44        def resolve_secret(env: Optional[dict] = None) -> Optional[str]:
45        """Resolve the SMTP password from env or GCP Secret Manager. Returns None if
unresolvable."""
46        env = env if env is not None else os.environ
47        direct = env.get("NIGHTLY_SMTP_PASS")
48        if direct:
49            return direct
50        secret_ref = env.get("NIGHTLY_SMTP_SECRET")
51        if not secret_ref:
52            return None
53        # Prefer the client lib; fall back to the gcloud CLI (no extra dependency required).
54        try:
55            from google.cloud import secretmanager # type: ignore
56            client = secretmanager.SecretManagerServiceClient()
57            return
client.access_secret_version(name=secret_ref).payload.data.decode("utf-8")

```

```

58         except Exception: # noqa: BLE001 ? lib absent or auth issue ? try the CLI
59             pass
60         try:
61             import subprocess
62             name = secret_ref.split("/secrets/", 1)[1].split("/", 1)[0] if "/secrets/" in
secret_ref else secret_ref
63             out = subprocess.run(["gcloud", "secrets", "versions", "access", "latest", "--
secret", name],
64                                   capture_output=True, text=True, timeout=30)
65             return out.stdout.strip() if out.returncode == 0 and out.stdout.strip() else
None
66         except Exception as e: # noqa: BLE001
67             log.warning("secret-manager resolution failed: %s", e)
68             return None
69
70
71     def send_report(subject: str, body_md: str, to: Optional[str] = None,
72                   env: Optional[dict] = None) -> bool:
73         """Send the report. Returns True if sent, False if skipped (no creds) ? never raises
on missing creds."""
74         env = env if env is not None else os.environ
75         to_addr = to or env.get("NIGHTLY_EMAIL_TO") or DEFAULT_TO
76         host = env.get("NIGHTLY_SMTP_HOST", DEFAULT_HOST)
77         port = int(env.get("NIGHTLY_SMTP_PORT", DEFAULT_PORT))
78         user = env.get("NIGHTLY_SMTP_USER") or to_addr
79         password = resolve_secret(env)
80         if not password:
81             log.warning("no SMTP password (set NIGHTLY_SMTP_PASS or NIGHTLY_SMTP_SECRET) ?
email skipped.")
82             return False
83
84         import smtplib
85         msg = build_message(subject, body_md, from_addr=user, to_addr=to_addr)
86         with smtplib.SMTP(host, port, timeout=60) as s:
87             s.starttls()
88             s.login(user, password)
89             s.send_message(msg)
90         log.info("emailed report to %s", to_addr)
91         return True
92

```

4. user (2026-06-22T09:47:36.841Z)

```

1     """Nightly END-TO-END engine regression ? reel-count + quality + cost (the GCP-VM
nightly).
2
3     Runs the **configured + checkpointed** pipeline (the real WASB GPU detector, as a
deployment serves
4     it) end-to-end on 2-3 short owned golden videos and asserts the **number of reels does
not change**
5     vs a frozen baseline. Unlike ``scripts/nightly_regression.py`` (PR #202 ? a CPU re-score
of CACHED
6     trajectories, the windowing/scorer only), this one runs the *full* pipeline on the *real
model*, so it
7     needs a GPU ? hence it runs on an ephemeral GCP GPU VM
8     (`deploy/gcp/nightly_regression/`), emails the
9     owner the report, and self-deletes. It is **default-additive**: a tripwire, not the
promotion gate.
10
11     What it asserts / reports, per the owner's ask ("quality metrics + regression report +
cost"):
12
13     ? REEL COUNT per video ? ``len(play_segments)`` after segmentation; the detected
highlight clips.
14
15     Pure-CV / WASB detection is deterministic for a fixed (video, config, weights), so
this is asserted

```

```

13         **EXACT** (with a re-run-once guard for the one known flake: a transient DB-insert-
None, see Sguard).
14     ? QUALITY METRICS per video (when golden labels are provided) ? the R2 tolerance-match
block from
15     ``backend.eval.rally_seg_eval`` (tol_f1 / f1@0.5 / f1@0.25), compared with a small
tolerance.
16     ? COST of the run ? VM uptime (or processing wall) × machine $/hr via
``backend.billing`` ? USD + INR.
17
18     Extensible pattern (add a video = one manifest row + re-freeze):
19     ``eval_baselines/reelcount_manifest.json`` ? the videos (gs:// / local / gdrive://
URIs) + segmenter
20     ``eval_baselines/reelcount_baseline.json`` ? the frozen expected counts (regenerated,
never hand-edited)
21
22     Usage
23     -----
24     Freeze the baseline from the current configured+checkpointed engine (after an intended
change, or first run):
25         python scripts/nightly_reelcount.py --update-baseline
26
27     Nightly check (exit 1 on regression; writes output/nightly_reelcount/<date>.md; emails
if --email):
28         python scripts/nightly_reelcount.py --email
29
30     Exit codes: 0 pass · 1 regression (reel-count changed / a video errored / quality
dropped) · 2 operational
31     (manifest/videos missing, nothing ran) · 3 no baseline · 4 stale (corpus/segmenter
changed ? re-freeze).
32     """
33     from __future__ import annotations
34
35     import argparse
36     import datetime as _dt
37     import json
38     import os
39     import sys
40     from dataclasses import dataclass, field
41     from typing import Any, Dict, List, Optional, Tuple
42
43     REPO_ROOT = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
44     if REPO_ROOT not in sys.path:
45         sys.path.insert(0, REPO_ROOT)
46
47     # billing is pure (typing only) ? safe to import at module top.
48     from backend import billing # noqa: E402
49
50     DEFAULT_MANIFEST = os.path.join(REPO_ROOT, "eval_baselines", "reelcount_manifest.json")
51     DEFAULT_BASELINE = os.path.join(REPO_ROOT, "eval_baselines", "reelcount_baseline.json")
52     DEFAULT_REPORT_DIR = os.path.join(REPO_ROOT, "output", "nightly_reelcount")
53
54     #: Quality metrics we report + (when labels exist) gate ? higher is better, small
tolerance.
55     QUALITY_HIGHER: Tuple[str, ...] = ("tol_f1", "f1@0.5", "f1@0.25")
56     DEFAULT_QUALITY_TOL = 0.02 # quality is float/labeled ? a looser tripwire than the
exact reel-count
57     DEFAULT_FX_INR_PER_USD = 83.0
58     #: Fallback machine price if the manifest doesn't pin one (T4 on n1-standard-8, asia-
south1, ~mid-2026).
59     DEFAULT_MACHINE_USD_PER_HR = 0.35
60
61
62     @dataclass(frozen=True)
63     class Tolerances:
64         quality_tol: float = DEFAULT_QUALITY_TOL

```

```

65
66     def to_dict(self) -> Dict[str, Any]:
67         return {"quality_tol": self.quality_tol}
68
69
70     # ----- #
71     # Pure cost math (wraps backend.billing ? no IO).
72     # ----- #
73     def cost_report(billed_seconds: float, machine_usd_per_hr: float,
74                     fx_inr_per_usd: float, gpu_seconds: float = 0.0) -> Dict[str, Any]:
75         """Run cost from billed wall-seconds × machine rate ? USD + INR (``backend.billing``
arithmetic).
76
77         ``billed_seconds`` is the VM uptime (boot?delete) when the orchestrator passes
``--vm-uptime-seconds``;
78         otherwise it's the in-process pipeline wall time (a lower bound ? the VM also pays
for boot/install)."""
79         rate_per_sec = float(machine_usd_per_hr) / 3600.0
80         usage = {billing.WALL_SECONDS: float(billed_seconds)}
81         rates = {billing.WALL_SECONDS: rate_per_sec}
82         usd, breakdown = billing.compute_cost_usd(usage, rates)
83         return {
84             "billed_seconds": round(float(billed_seconds), 1),
85             "gpu_seconds": round(float(gpu_seconds), 1),
86             "machine_usd_per_hr": float(machine_usd_per_hr),
87             "usd": usd,
88             "inr": billing.to_inr(usd, fx_inr_per_usd),
89             "breakdown": breakdown,
90         }
91
92
93     # ----- #
94     # Pure summary + comparison core (no IO ? unit-testable with plain dicts).
95     # ----- #
96     def summarize_results(run_results: List[Dict[str, Any]], segmenter: str,
97                           cost: Dict[str, Any]) -> Dict[str, Any]:
98         """Compact, comparable snapshot from per-video run results.
99
100         Each run result: ``{name, reel_count, ok, error, quality: {tol_f1,...}|None,
wall_seconds}``.
101         A video that errored keeps ``reel_count=None`` + ``ok=False`` (so the diff flags it,
never hides it)."""
102         per_video: Dict[str, Any] = {}
103         for r in run_results:
104             per_video[r["name"]] = {
105                 "reel_count": r.get("reel_count"),
106                 "ok": bool(r.get("ok", False)),
107                 "error": r.get("error"),
108                 "quality": r.get("quality"),
109             }
110         return {"segmenter": segmenter, "cost": cost, "per_video": per_video,
111               "n": len(per_video)}
112
113
114     @dataclass
115     class Finding:
116         video: str
117         metric: str # "reel_count" | "error" | a quality key | "segmenter" |
"corpus"
118         kind: str # "regression" | "improvement" | "stale"
119         baseline: Optional[float] = None
120         current: Optional[float] = None
121         detail: str = ""
122
123         def message(self) -> str:

```

```

124         if self.kind == "stale":
125             return f"[STALE] {self.metric}: {self.detail}"
126         if self.metric == "error":
127             return f"[REGRESSION] {self.video}: pipeline ERROR ? {self.detail}"
128         b = "?" if self.baseline is None else f"{self.baseline:g}"
129         c = "?" if self.current is None else f"{self.current:g}"
130         tag = "REGRESSION" if self.kind == "regression" else "improvement"
131         return f"[{tag}] {self.video}/{self.metric}: {b} ? {c}"
132
133
134     @dataclass
135     class NightlyResult:
136         findings: List[Finding] = field(default_factory=list)
137         added: List[str] = field(default_factory=list)
138         removed: List[str] = field(default_factory=list)
139
140         @property
141         def regressions(self) -> List[Finding]:
142             return [f for f in self.findings if f.kind == "regression"]
143
144         @property
145         def improvements(self) -> List[Finding]:
146             return [f for f in self.findings if f.kind == "improvement"]
147
148         @property
149         def stale(self) -> List[Finding]:
150             return [f for f in self.findings if f.kind == "stale"]
151
152         def overall_status(self) -> str:
153             if self.regressions:
154                 return "REGRESSION"
155             if self.stale:
156                 return "STALE"
157             return "PASS"
158
159         def exit_code(self) -> int:
160             if self.regressions:
161                 return 1
162             if self.stale:
163                 return 4
164             return 0
165
166
167     def compare(baseline: Dict[str, Any], current: Dict[str, Any], tols: Tolerances) ->
NightlyResult:
168         """Diff current run vs frozen baseline.
169
170         * Segmenter changed ? STALE (numbers not comparable; re-freeze).
171         * Corpus (video set) changed ? STALE + report added/removed; still diff the
INTERSECTION.
172         * Per shared video: reel_count EXACT (any change = regression); a pipeline ERROR =
regression;
173         quality metrics drop beyond ``quality_tol`` = regression (rise = improvement).
174         """
175         res = NightlyResult()
176         if baseline.get("segmenter") != current.get("segmenter"):
177             res.findings.append(Finding("", "segmenter", "stale",
178                                     detail=f"{baseline.get('segmenter')} ?
{current.get('segmenter')}"))
179             return res
180
181         b_pv, c_pv = baseline.get("per_video", {}), current.get("per_video", {})
182         res.added = sorted(set(c_pv) - set(b_pv))
183         res.removed = sorted(set(b_pv) - set(c_pv))
184         if res.added or res.removed:

```

```

185         res.findings.append(Finding("", "corpus", "stale",
186                                 detail=f"added={res.added} removed={res.removed}"))
187
188     for v in sorted(set(b_pv) & set(c_pv)):
189         b, c = b_pv[v], c_pv[v]
190         # 1) pipeline must still succeed
191         if not c.get("ok", False) or c.get("reel_count") is None:
192             res.findings.append(Finding(v, "error", "regression",
193                                     detail=str(c.get("error") or "no reel_count
194 produced")))
195             continue
196         # 2) reel count ? EXACT
197         bc, cc = b.get("reel_count"), c.get("reel_count")
198         if bc is not None and cc is not None and cc != bc:
199             res.findings.append(Finding(v, "reel_count", "regression", float(bc),
200 float(cc)))
201         # 3) quality metrics ? tolerance (only if both sides have them)
202         bq, cq = b.get("quality") or {}, c.get("quality") or {}
203         for m in QUALITY_HIGHER:
204             bv, cv = bq.get(m), cq.get(m)
205             if bv is None or cv is None:
206                 continue
207             if cv < bv - tols.quality_tol:
208                 res.findings.append(Finding(v, m, "regression", bv, cv))
209             elif cv > bv + tols.quality_tol:
210                 res.findings.append(Finding(v, m, "improvement", bv, cv))
211
212     return res
213
214 # ----- #
215 # Report formatting (pure).
216 # ----- #
217
218 def _q(qd: Optional[Dict[str, Any]], key: str) -> str:
219     if not qd or qd.get(key) is None:
220         return "?"
221     return f"{qd[key]:.3f}"
222
223
224 def format_report(baseline: Dict[str, Any], current: Dict[str, Any],
225                  result: NightlyResult, date: str, head: str) -> str:
226     status = result.overall_status()
227     badge = {"PASS": "? PASS", "REGRESSION": "? REGRESSION",
228             "STALE": "? STALE (re-freeze the baseline)"}[status]
229     cost = current.get("cost", {})
230     L: List[str] = [
231         f"# Nightly engine regression (reel-count + quality + cost) ? {date}",
232         "",
233         f"***Status: {badge}** · exit code {result.exit_code()} · segmenter
234 {current.get('segmenter')}",
235         "",
236         f"- HEAD {head} · baseline {baseline.get('git_commit', '?')} (frozen
237 {baseline.get('generated_at', '?')})",
238         f"- **Run cost: ${cost.get('usd', 0):.4f} (?{cost.get('inr', 0):.2f})** · "
239         f"billed {cost.get('billed_seconds', 0):.0f}s @ ${cost.get('machine_usd_per_hr',
240 0)}/hr",
241         "",
242     ]
243     if result.regressions:
244         L += ["## ? Regressions", ""] + [f"- {f.message()}" for f in result.regressions]
245     + [""]
246     if result.stale:
247         L += ["## ? Stale / not comparable", ""] + [f"- {f.message()}" for f in
248 result.stale]
249     L += ["- Action: `python scripts/nightly_reelcount.py --update-baseline` (after
250 confirming the change is intended).", ""]

```

```

242
243 # Per-video table.
244 b_pv, c_pv = baseline.get("per_video", {}), current.get("per_video", {})
245 L += ["## Per-video", "",
246       "| video | reels (base?now) | tol_f1 (base?now) | f1@0.5 | status |",
247       "|---|---|---|---|---|"]
248 for v in sorted(set(b_pv) | set(c_pv)):
249     b, c = b_pv.get(v, {}), c_pv.get(v, {})
250     bc = b.get("reel_count")
251     cc = c.get("reel_count")
252     rc = f"'?' if bc is None else bc?{'ERROR' if not c.get('ok', True) else ('?'
if cc is None else cc)}"
253     bq, cq = b.get("quality"), c.get("quality")
254     tf = f"_{q(bq, 'tol_f1')}_?_{q(cq, 'tol_f1')}"
255     f5 = f"_{q(bq, 'f1@0.5')}_?_{q(cq, 'f1@0.5')}"
256     vstatus = "?" if any(f.video == v and f.kind == "regression" for f in
result.findings) else ""
257     L.append(f"| {v} | {rc} | {tf} | {f5} | {vstatus} |")
258     L.append("")
259
260 if result.improvements:
261     L += ["_Improvements over baseline (re-freeze to ratchet up):_" \
262          + [f"- {f.message()}" for f in result.improvements] + [""]
263     L += ["---",
264           "_Full configured+checkpointed pipeline on the real model (GCP GPU VM).
Tripwire only ? does "
265           "not change the promotion gate. Complements PR #202's cached-trajectory CPU
re-score._"]
266     return "\n".join(L)
267
268
269 # ----- #
270 # IO shell ? running the configured pipeline per video (needs backend + GPU + video).
271 # ----- #
272 def _load_gts(labels_ref: Optional[str]) -> Optional[List[Tuple[float, float]]]:
273     """Load ground-truth rally intervals from a golden-features JSON (`gts` key) or a
rallies CSV
274     (`start,end` columns / pairs). Returns None when no labels are configured."""
275     if not labels_ref:
276         return None
277     from backend.storage import fetch
278     local = fetch(labels_ref)
279     if local.endswith(".json"):
280         with open(local, encoding="utf-8") as fh:
281             data = json.load(fh)
282             return [(float(s), float(e)) for s, e in (data.get("gts") or [])]
283     # CSV: header start,end (or first two numeric columns)
284     import csv
285     out: List[Tuple[float, float]] = []
286     with open(local, encoding="utf-8") as fh:
287         for row in csv.reader(fh):
288             try:
289                 out.append((float(row[0]), float(row[1])))
290             except (ValueError, IndexError):
291                 continue
292     return out
293
294
295 def _segments_to_intervals(segs: List[Dict[str, Any]]) -> List[Tuple[float, float]]:
296     """(`start, end` pairs from segmenter result dicts, tolerant of the two key
conventions in the
297     codebase (`start_time`/`end_time` ? motion/DB shape ? or `start`/`end`).
Used only for the
298     optional quality metrics; the reel COUNT is `len(segs)` regardless, so a key
mismatch degrades

```

```

299     quality scoring gracefully (skipped) without ever miscounting reels."""
300     out: List[Tuple[float, float]] = []
301     for s in segs:
302         start = s.get("start_time", s.get("start"))
303         end = s.get("end_time", s.get("end"))
304         if start is not None and end is not None:
305             out.append((float(start), float(end)))
306     return out
307
308
309     def run_one_video(video: Dict[str, Any], segmenter: str, config: Any,
310                     rerun_on_mismatch: bool = True,
311                     baseline_count: Optional[int] = None) -> Dict[str, Any]:
312         """Run the configured pipeline on one video ? reel count (+ quality vs labels). Lazy
backend imports.
313
314         The re-run-once guard (the one known non-determinism ? a transient
``add_play_segment``?None drop,
315         database.py): if the count != the frozen baseline, re-process the video ONCE; a
transient drop won't
316         reproduce, a real change will."""
317         import random
318         import time
319
320         from backend.eval import rally_seg_eval
321         from backend.infrastructure.database import Database
322         from backend.pipeline.segmenters import segmenter_registry
323         from backend.results import Failure
324         from backend.storage import fetch
325         from backend.utils.hashing import compute_video_id
326
327         name = video["name"]
328         try:
329             local = fetch(video["uri"])
330         except Exception as e: # noqa: BLE001 ? surface, never hide
331             return {"name": name, "reel_count": None, "ok": False, "error": f"fetch failed:
{e}",
332                     "quality": None, "wall_seconds": 0.0}
333
334         gts = _load_gts(video.get("labels_uri"))
335         seg_cls = segmenter_registry.get(segmenter)
336         if not seg_cls:
337             return {"name": name, "reel_count": None, "ok": False,
338                     "error": f"segmenter '{segmenter}' not registered", "quality": None,
"wall_seconds": 0.0}
339
340         def _process(db_path: str) -> Tuple[Optional[int], Optional[List[Tuple[float,
float]]], float, Optional[str]]:
341             random.seed(0) # close motion.py's unseeded metadata RNG so secondary fields
are reproducible too
342             db = Database(db_path)
343             video_id = compute_video_id(local)
344             t0 = time.monotonic()
345             result = seg_cls(db, config).process_video(video_id, local)
346             wall = time.monotonic() - t0
347             if isinstance(result, Failure):
348                 return None, None, wall, getattr(result, "message", "segmenter failed")
349             segs = result.value if hasattr(result, "value") else result
350             intervals = _segments_to_intervals(segs)
351             return len(segs), intervals, wall, None
352
353         import tempfile
354         total_wall = 0.0
355         with tempfile.TemporaryDirectory(prefix="nightly_reel_") as td:
356             count, intervals, wall, err = _process(os.path.join(td, "run1.db"))

```

```

357         total_wall += wall
358     if err is not None:
359         return {"name": name, "reel_count": None, "ok": False, "error": err,
360                "quality": None, "wall_seconds": round(total_wall, 2)}
361     # re-run-once guard for the transient DB-drop flake
362     if rerun_on_mismatch and baseline_count is not None and count != baseline_count:
363         count2, intervals2, wall2, err2 = _process(os.path.join(td, "run2.db"))
364         total_wall += wall2
365         if err2 is None and count2 is not None:
366             count, intervals = count2, intervals2 # the reproduced (or corrected)
value
367
368         quality = None
369         if gts and intervals is not None:
370             row = rally_seg_eval.score_intervals(intervals, gts)
371             quality = {k: row[k] for k in ("tol_f1", "f1@0.5", "f1@0.25",
"tol_precision", "tol_recall")}
372                 if k in row}
373
374     return {"name": name, "reel_count": count, "ok": True, "error": None,
375            "quality": quality, "wall_seconds": round(total_wall, 2)}
376
377
378     def load_manifest(path: str) -> Dict[str, Any]:
379         with open(path, encoding="utf-8") as fh:
380             return json.load(fh)
381
382
383     def run_corpus(manifest: Dict[str, Any], baseline: Optional[Dict[str, Any]],
384                  rerun_on_mismatch: bool = True) -> Tuple[List[Dict[str, Any]], float]:
385         """Run every manifest video through the configured pipeline. Returns (run_results,
total_wall_s)."""
386         from backend.config import load_config
387
388         config = load_config()
389         # apply manifest config overrides onto the typed config (dotted keys)
390         for dotted, val in (manifest.get("config_overrides") or {}).items():
391             _set_dotted(config, dotted, val)
392         segmenter = manifest.get("segmenter") or getattr(getattr(config, "indexing", None),
"default_segmenter", "motion")
393         b_pv = (baseline or {}).get("per_video", {})
394         results: List[Dict[str, Any]] = []
395         total_wall = 0.0
396         for v in manifest.get("videos", []):
397             base_count = (b_pv.get(v["name"]) or {}).get("reel_count")
398             r = run_one_video(v, segmenter, config, rerun_on_mismatch, base_count)
399             total_wall += float(r.get("wall_seconds") or 0.0)
400             results.append(r)
401         return results, total_wall
402
403
404     def _set_dotted(config: Any, dotted: str, value: Any) -> None:
405         """Best-effort dotted-path set onto a pydantic config (e.g.
'indexing.motion_threshold')."""
406         parts = dotted.split(".")
407         obj = config
408         for p in parts[:-1]:
409             obj = getattr(obj, p, None)
410             if obj is None:
411                 return
412         try:
413             setattr(obj, parts[-1], value)
414         except Exception: # noqa: BLE001 ? overrides are best-effort; a bad key shouldn't
crash the run
415             pass

```

```

416
417
418 def _git_head() -> str:
419     import subprocess
420     try:
421         out = subprocess.run(["git", "-C", REPO_ROOT, "rev-parse", "--short", "HEAD"],
422                               capture_output=True, text=True, timeout=10)
423         return out.stdout.strip() or "?"
424     except Exception: # noqa: BLE001
425         return "?"
426
427
428 def load_baseline(path: str) -> Optional[Dict[str, Any]]:
429     if not os.path.exists(path):
430         return None
431     with open(path, encoding="utf-8") as fh:
432         return json.load(fh)
433
434
435 def save_baseline(path: str, snapshot: Dict[str, Any]) -> None:
436     out = dict(snapshot)
437     out["generated_at"] = _dt.date.today().isoformat()
438     out["git_commit"] = _git_head()
439     out["_doc"] = ("Frozen nightly reel-count baseline. Regenerate with `python
scripts/nightly_reelcount.py "
440                   "--update-baseline` after an INTENDED engine change or when adding a
video. Tied to the "
441                   "configured+checkpointed engine + the manifest corpus.")
442     out.pop("cost", None) # cost is per-run, not part of the frozen baseline
443     os.makedirs(os.path.dirname(path), exist_ok=True)
444     tmp = path + ".tmp"
445     with open(tmp, "w", encoding="utf-8") as fh:
446         json.dump(out, fh, indent=2, sort_keys=True)
447         fh.write("\n")
448     os.replace(tmp, path)
449
450
451 def build_parser() -> argparse.ArgumentParser:
452     p = argparse.ArgumentParser(description="Nightly E2E engine regression: reel-count +
quality + cost.")
453     p.add_argument("--manifest", default=DEFAULT_MANIFEST)
454     p.add_argument("--baseline", default=DEFAULT_BASELINE)
455     p.add_argument("--report-dir", default=DEFAULT_REPORT_DIR)
456     p.add_argument("--quality-tol", type=float, default=DEFAULT_QUALITY_TOL)
457     p.add_argument("--vm-uptime-seconds", type=float, default=None,
458                   help="authoritative VM uptime (boot?delete) for the cost line; "
459                   "defaults to the in-process pipeline wall time (a lower bound)")
460     p.add_argument("--no-rerun-guard", action="store_true",
461                   help="disable the re-run-once-on-mismatch guard (the DB-drop flake
mitigation)")
462     p.add_argument("--update-baseline", action="store_true",
463                   help="freeze the current reel counts + quality as the new baseline
instead of checking")
464     p.add_argument("--email", action="store_true", help="email the report (see
scripts/nightly_email.py)")
465     p.add_argument("--email-to", default=None, help="override the report recipient")
466     p.add_argument("--date", default=None, help="override the report date (YYYY-MM-DD;
for tests)")
467     p.add_argument("--no-report", action="store_true", help="skip writing the dated
report file")
468     return p
469
470
471 def main(argv: Optional[List[str]] = None) -> int:
472     import logging

```

```

473 logging.basicConfig(level=logging.INFO, format="%(levelname)s %(message)s")
474 log = logging.getLogger("nightly_reelcount")
475 args = build_parser().parse_args(argv)
476 tols = Tolerances(quality_tol=args.quality_tol)
477 date = args.date or _dt.date.today().isoformat()
478
479 if not os.path.exists(args.manifest):
480     log.error("manifest not found: %s", args.manifest)
481     return 2
482 manifest = load_manifest(args.manifest)
483 baseline = load_baseline(args.baseline)
484
485 try:
486     run_results, total_wall = run_corpus(manifest, baseline, rerun_on_mismatch=not
args.no_rerun_guard)
487 except Exception as e: # noqa: BLE001 ? operational failure (deps/config) ? exit 2,
never a false "pass"
488     log.error("run failed: %s", e)
489     return 2
490 if not run_results:
491     log.error("no videos in manifest %s", args.manifest)
492     return 2
493
494 billed = args.vm_uptime_seconds if args.vm_uptime_seconds is not None else
total_wall
495 cost = cost_report(
496     billed_seconds=billed,
497     machine_usd_per_hr=float(manifest.get("machine_usd_per_hr",
DEFAULT_MACHINE_USD_PER_HR)),
498     fx_inr_per_usd=float(manifest.get("fx_inr_per_usd", DEFAULT_FX_INR_PER_USD)),
499     gpu_seconds=total_wall,
500 )
501 segmenter = manifest.get("segmenter", "motion")
502 snapshot = summarize_results(run_results, segmenter, cost)
503 snapshot["git_commit"] = _git_head()
504 snapshot["generated_at"] = date
505
506 if args.update_baseline:
507     save_baseline(args.baseline, snapshot)
508     log.info("baseline frozen ? %s", args.baseline)
509     for r in run_results:
510         log.info(" %s: reels=%s ok=%s", r["name"], r.get("reel_count"),
r.get("ok"))
511     return 0
512
513 if baseline is None:
514     log.error("no baseline at %s ? freeze one with --update-baseline (on the
configured engine).", args.baseline)
515     return 3
516
517 result = compare(baseline, snapshot, tols)
518 report = format_report(baseline, snapshot, result, date, snapshot["git_commit"])
519
520 if not args.no_report:
521     os.makedirs(args.report_dir, exist_ok=True)
522     md_path = os.path.join(args.report_dir, f"{date}.md")
523     with open(md_path, "w", encoding="utf-8") as fh:
524         fh.write(report + "\n")
525     with open(os.path.join(args.report_dir, f"{date}.json"), "w", encoding="utf-8")
as fh:
526         json.dump({"date": date, "status": result.overall_status(), "exit_code":
result.exit_code(),
527                   "snapshot": snapshot, "regressions": [f.message() for f in
result.regressions]},
528                   fh, indent=2, sort_keys=True)

```

```

529         log.info("report ? %s", md_path)
530
531     print(report)
532
533     if args.email:
534         try:
535             from scripts.nightly_email import send_report
536             subject = f"[KhelSutra nightly] {result.overall_status()} ? reel-count
engine regression {date}"
537             sent = send_report(subject, report, to=args.email_to)
538             log.info("email %s", "sent" if sent else "skipped (no creds ? report still
written)")
539         except Exception as e: # noqa: BLE001 ? a broken mailer must NOT mask the
engine regression verdict
540             log.error("email failed (report still written): %s", e)
541
542     code = result.exit_code()
543     if code == 1:
544         log.error("REEL-COUNT REGRESSION (%d) ? see %s", len(result.regressions),
args.report_dir)
545     elif code == 4:
546         log.warning("baseline STALE ? corpus/segmenter changed; re-freeze with --update-
baseline")
547     else:
548         log.info("nightly reel-count check PASSED (cost $%.4f / ?%.2f)", cost["usd"],
cost["inr"])
549     return code
550
551
552 if __name__ == "__main__":
553     raise SystemExit(main())
554

```

5. user (2026-06-22T09:47:37.408Z)

```

1     #!/usr/bin/env bash
2     # Startup-script for the NIGHTLY engine-regression VM. Runs UNATTENDED on boot, then
DELETES the VM
3     # (cost rule: a running GPU VM is the only real cost). Survives SSH drops (it's a
startup-script, not
4     # an interactive session). Reads its config from instance metadata (set by launch.sh).
5     #
6     # Flow: Ops Agent ? pull repo @ HEAD from GCS ? build venv + WASB env ? stage
weights ?
7     #         run scripts/nightly_reelcount.py on the gs:// golden clips (emails the report)
?
8     #         upload report to GCS ? SELF-DELETE (even on failure, via the EXIT trap).
9     #
10    # Metadata keys (launch.sh sets these): repo-tgz-uri, git-sha, gcs-bucket, report-
prefix,
11    # weights-uri, email (true/false). NOT verified in CI ? dry-run via launch.sh on a
real VM first.
12    set -uo pipefail
13    exec > >(tee -a /var/log/nightly-regression.log) 2>&1
14    echo "==== nightly-regression startup $(date -u +%FT%TZ) ====="
15    BOOT_EPOCH=$(date +%s)
16
17    md() { curl -s -H "Metadata-Flavor: Google"
"http://metadata.google.internal/computeMetadata/v1/$1"; }
18    NAME=$(md instance/name)
19    ZONE=$(md instance/zone | awk -F/ '{print $NF}')
20    REPO_TGZ=$(md instance/attributes/repo-tgz-uri || true)
21    GIT_SHA=$(md instance/attributes/git-sha || echo unknown)
22    BUCKET=$(md instance/attributes/gcs-bucket || true)
23    REPORT_PREFIX=$(md instance/attributes/report-prefix || echo nightly/reports)

```

```

24     WEIGHTS_URI=$(md instance/attributes/weights-uri || true)
25     DO_EMAIL=$(md instance/attributes/email || echo true)
26     SMTP_SECRET=$(md instance/attributes/smtp-secret || true)
27     SMTP_USER=$(md instance/attributes/smtp-user || true)
28     EMAIL_TO=$(md instance/attributes/email-to || echo avi.dullu@gmail.com)
29     DATE=$(date -u +%F)
30
31     # ALWAYS delete the VM on exit (success OR failure) ? never leave a billing GPU VM
behind.
32     cleanup() {
33         echo "==== self-delete $NAME ($ZONE) ====="
34         gcloud compute instances delete "$NAME" --zone="$ZONE" --quiet || \
35             echo "WARN: self-delete failed ? DELETE MANUALLY: gcloud compute instances delete
$NAME --zone=$ZONE -q"
36     }
37     trap cleanup EXIT
38
39     fail_email() { # best-effort failure alert (the report email won't have fired on an
early crash)
40         echo "RUN FAILED ? attempting a failure alert"
41         if [ "$DO_EMAIL" = "true" ] && [ -d ~/rally ]; then
42             cd ~/rally 2>/dev/null && . .venv/bin/activate 2>/dev/null && \
43                 NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
44                 python - <<PY || true
45         from scripts.nightly_email import send_report
46         send_report("[KhelSutra nightly] ERROR ? engine regression VM run failed ($DATE)",
47                     "The nightly engine-regression run failed on the VM before producing a
report.\n"
48                     "See gs://$BUCKET/$REPORT_PREFIX/$DATE/ and the VM serial
log.\nGIT_SHA=$GIT_SHA", to="$EMAIL_TO")
49         PY
50         fi
51     }
52     trap 'fail_email' ERR
53
54     # 1) Observability (Ops Agent) ? same as create_gpu_vm.sh bakes in.
55     curl -sSO https://dl.google.com/cloudagents/add-google-cloud-ops-agent-repo.sh && \
56         sudo bash add-google-cloud-ops-agent-repo.sh --also-install || echo "WARN: ops-agent
install failed (non-fatal)"
57
58     # 2) Pull the repo @ HEAD (staged to GCS by launch.sh ? no GitHub auth needed; uses the
worker SA).
59     mkdir -p ~/rally && cd ~/rally
60     gcloud storage cp "$REPO_TGZ" /tmp/rally.tgz
61     tar -xzf /tmp/rally.tgz -C ~/rally
62     echo "$GIT_SHA" > ~/rally/GIT_SHA
63
64     # 3) Build the env + native WASB (provider-agnostic DO scripts, reused on GCP per
deploy/gcp/README §4).
65     sudo bash deploy/digitalocean/mount_scratch.sh 2>/dev/null || true
66     export TMPDIR=/scratch 2>/dev/null || true
67     ./deploy/digitalocean/bootstrap.sh --gpu
68     . .venv/bin/activate
69     pip install -e '[storage-gcs]'
70     bash deploy/digitalocean/gpu_setup.sh || echo "WARN: gpu_setup (WASB env) returned
nonzero ? continuing"
71
72     # 4) Stage the WASB checkpoint (gpu_setup does NOT fetch weights ? README §4).
73     if [ -n "$WEIGHTS_URI" ]; then
74         mkdir -p ~/models/WASB-SBDT/pretrained_weights
75         gcloud storage cp "$WEIGHTS_URI" ~/models/WASB-
SBDT/pretrained_weights/wasb_badminton_best.pth.tar
76         export WASB_WEIGHTS=~/models/WASB-SBDT/pretrained_weights/wasb_badminton_best.pth.tar
77         fi

```

```

78
79     # 5) Run the regression on the gs:// golden clips in the committed manifest. Cost = VM
uptime so far.
80     UPTIME=$(( (date +%s) - BOOT_EPOCH ))
81     EMAIL_FLAG=""; [ "$DO_EMAIL" = "true" ] && EMAIL_FLAG="--email --email-to $EMAIL_TO"
82     set +e
83     NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
84     python scripts/nightly_reelcount.py --vm-uptime-seconds "$UPTIME" $EMAIL_FLAG
85     RUN_RC=$?
86     set -e
87     echo "runner exit code: $RUN_RC (0 pass / 1 regression / 4 stale)"
88
89     # 6) Upload the dated report to GCS for the record (alongside the email).
90     if [ -n "$BUCKET" ] && [ -d output/nightly_reelcount ]; then
91         gcloud storage cp -r "output/nightly_reelcount/$DATE."*
"gs://$BUCKET/$REPORT_PREFIX/$DATE/" || \
92         gcloud storage cp -r output/nightly_reelcount "gs://$BUCKET/$REPORT_PREFIX/$DATE/"
|| true
93     fi
94
95     echo "==== nightly-regression done rc=$RUN_RC $(date -u +%FT%TZ) ====="
96     trap - ERR # disable the failure-email trap; the EXIT (self-delete) trap stays
97     # the EXIT trap deletes the VM
98

```

6. user (2026-06-22T09:47:37.729Z)

```

1     #!/usr/bin/env bash
2     # Launch ONE nightly engine-regression run on an ephemeral GCP GPU VM, then return. The
VM runs
3     # run_on_vm.sh as its startup-script (regression ? email ? SELF-DELETE), so this just
stages the repo
4     # and kicks the VM off. Driver-agnostic: call it from Cloud Scheduler?Cloud Run job, a
cron on any
5     # always-on box, the owner's machine, or a Claude routine. Needs gcloud auth + the repo
checkout.
6     #
7     # bash deploy/gcp/nightly_regression/launch.sh # L4, stages HEAD,
kicks the VM
8     # RALLY_GPU=t4 RALLY_REF=$(git rev-parse HEAD) bash
deploy/gcp/nightly_regression/launch.sh
9     #
10    # ? Run this MANUALLY once as a dry-run (watch the serial log) before wiring a scheduler
? the VM-side
11    # flow is NOT covered by CI. Tail: gcloud compute instances get-serial-port-output
<name> --zone=<z>.
12    set -uo pipefail
13
14    PROJECT="${RALLY_GCP_PROJECT:-gen-lang-client-0306026826}"
15    BUCKET="${RALLY_BUCKET:-kheldutra-rally-corpus}"
16    SA="rally-worker@${PROJECT}.iam.gserviceaccount.com"
17    GPU="${RALLY_GPU:-l4}"
18    NAME="${RALLY_VM_NAME:-rally-nightly-$(date -u +%Y%m%d)}"
19    REF="${RALLY_REF:-HEAD}"
20    WEIGHTS_URI="${RALLY_WEIGHTS_URI:-gs://$BUCKET/weights/wasb_badminton_best.pth.tar}"
21    REPORT_PREFIX="${RALLY_REPORT_PREFIX:-nightly/reports}"
22    DO_EMAIL="${RALLY_EMAIL:-true}"
23    SMTP_SECRET="${NIGHTLY_SMTP_SECRET:-projects/$PROJECT/secrets/nightly-smtp-
pass/versions/latest}"
24    SMTP_USER="${NIGHTLY_SMTP_USER:-}"
25    EMAIL_TO="${NIGHTLY_EMAIL_TO:-avi.dullu@gmail.com}"
26    HERE="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
27    STARTUP="$HERE/run_on_vm.sh"
28    export CLOUDSDK_CORE_DISABLE_PROMPTS=1

```

```

29
30     ZONES="${RALLY_GCP_ZONES:-asia-south1-a asia-south1-b asia-south1-c asia-southeast1-a
asia-southeast1-b asia-southeast1-c}"
31     if [ "$GPU" = "l4" ]; then MACHINE="g2-standard-4"; ACCEL=(); else
MACHINE="n1-standard-8"; ACCEL=(--accelerator=type=nvidia-tesla-t4,count=1); fi
32
33     # 1) Stage the repo @ REF to GCS (no GitHub auth on the VM; the worker SA can read the
bucket).
34     SHA="$(git rev-parse "$REF")"
35     TGZ_URI="gs://$BUCKET/nightly/repo/${SHA}.tgz"
36     echo "== staging repo $SHA -> $TGZ_URI =="
37     TMP_TGZ="$(mktemp --suffix=.tgz)"
38     git archive --format=tar.gz -o "$TMP_TGZ" "$REF"
39     gcloud storage cp "$TMP_TGZ" "$TGZ_URI" --project="$PROJECT"
40     rm -f "$TMP_TGZ"
41
42     # 2) Zone-sweep create the VM (mirrors create_gpu_vm.sh) with run_on_vm.sh + the config
metadata.
43     ERR="$(mktemp)"
44     for Z in $ZONES; do
45         echo "== creating $NAME ($GPU) in $Z =="
46         if gcloud compute instances create "$NAME" --project="$PROJECT" --zone="$Z" \
47             --machine-type="$MACHINE" "${ACCEL[@]}" \
48             --maintenance-policy=TERMINATE --provisioning-model=STANDARD \
49             --image-family=common-cu129-ubuntu-2204-nvidia-580 --image-project=deeplearning-
platform-release \
50             --boot-disk-size=120GB --boot-disk-type=pd-balanced \
51             --service-account="$SA" --scopes=cloud-platform \
52             --metadata-from-file=startup-script="$STARTUP" \
53             --metadata=repo-tgz-uri="$TGZ_URI",git-sha="$SHA",gcs-bucket="$BUCKET",report-
prefix="$REPORT_PREFIX",weights-uri="$WEIGHTS_URI",email="$DO_EMAIL",smtp-
secret="$SMTP_SECRET",smtp-user="$SMTP_USER",email-to="$EMAIL_TO" \
54             2>"$ERR"; then
55             echo "? launched $NAME in $Z ? it will run the regression, email $EMAIL_TO, and
self-delete."
56             echo "    serial log: gcloud compute instances get-serial-port-output $NAME --zone=$Z
--project=$PROJECT"
57             echo "    if it doesn't self-delete: gcloud compute instances delete $NAME --zone=$Z
-q"
58             rm -f "$ERR"; exit 0
59         fi
60         if grep -qiE "STOCKOUT|does not have enough" "$ERR"; then echo "    (no $GPU capacity in
$Z ? next zone)"; continue; fi
61         echo "    create failed:"; tail -4 "$ERR"; rm -f "$ERR"; exit 1
62     done
63     echo "? no $GPU capacity in any swept zone ($ZONES). Retry later or RALLY_GPU=t4."
64     rm -f "$ERR"; exit 1
65

```

7. assistant (2026-06-22T09:47:40.605Z)

Now let me continue reading the remaining files:

8. user (2026-06-22T09:47:41.846Z)

```

1     #!/usr/bin/env bash
2     # Wire the nightly engine-regression as a CLOUD-NATIVE, scale-to-zero job:
3     #   Cloud Scheduler (cron) ? Cloud Run JOB (the launcher image) ? launch.sh
4     #   ? ephemeral GPU VM (regression + email + SELF-DELETE).
5     # The only always-on cost is $0 (Cloud Run jobs + Scheduler are scale-to-zero; the GPU
VM is ephemeral).
6     #
7     # ? OWNER-RUN, ONCE, after a successful manual `bash
deploy/gcp/nightly_regression/launch.sh` dry-run.
8     # NOT verified in CI (no GCP/GPU here). Prereqs: GPU quota (README S2), the golden clips

```

```

+ weights in
9      # gs://<bucket>, and the SMTP secret (see README "Enable").
10     #
11     #   bash deploy/gcp/nightly_regression/register_scheduler.sh           # build + deploy
+ schedule (03:00 IST)
12     #   RALLY_CRON='0 3 * * *' RALLY_TZ='Asia/Kolkata' bash ../register_scheduler.sh
13     set -euo pipefail
14
15     PROJECT="${RALLY_GCP_PROJECT:-gen-lang-client-0306026826}"
16     REGION="${RALLY_RUN_REGION:-asia-south1}"
17     JOB="${RALLY_RUN_JOB:-nightly-regression-launcher}"
18     SCHED="${RALLY_SCHED_NAME:-nightly-regression}"
19     CRON="${RALLY_CRON:-0 3 * * *}"           # 03:00 daily
20     TZ="${RALLY_TZ:-Asia/Kolkata}"
21     SA="rally-worker@${PROJECT}.iam.gserviceaccount.com"
22     IMAGE="${RALLY_LAUNCHER_IMAGE:-${REGION}-docker.pkg.dev/${PROJECT}/rally/nightly-
launcher:latest}"
23     REPO_ROOT="$(git rev-parse --show-toplevel)"
24     export CLOUDSDK_CORE_DISABLE_PROMPTS=1
25
26     echo "== enable APIs (run, scheduler, build, artifact registry) =="
27     gcloud services enable run.googleapis.com cloudscheduler.googleapis.com \
28         cloudbuild.googleapis.com artifactregistry.googleapis.com --project="$PROJECT"
29
30     echo "== ensure the Artifact Registry repo 'rally' exists =="
31     gcloud artifacts repositories describe rally --location="$REGION" --project="$PROJECT"
>/dev/null 2>&1 || \
32     gcloud artifacts repositories create rally --repository-format=docker
--location="$REGION" --project="$PROJECT"
33
34     echo "== build + push the launcher image from the repo root =="
35     ( cd "$REPO_ROOT" && gcloud builds submit --project="$PROJECT" --tag "$IMAGE" \
36         --gcs-source-staging-dir="gs://${RALLY_BUILD_BUCKET:-khelsutra-rally-
corpus}/cloudbuild" \
37         -f deploy/gcp/nightly_regression/Dockerfile.launcher . )
38
39     echo "== create/update the Cloud Run JOB (launcher needs compute-admin to create the GPU
VM) =="
40     gcloud run jobs deploy "$JOB" --project="$PROJECT" --region="$REGION" \
41         --image="$IMAGE" --service-account="$SA" --max-retries=0 --task-timeout=900s \
42         --set-env-vars="RALLY_GPU=${RALLY_GPU:-
l4},RALLY_EMAIL=true,NIGHTLY_EMAIL_TO=${NIGHTLY_EMAIL_TO:-avi.dullu@gmail.com}"
43
44     echo "== schedule it (Cloud Scheduler ? Run job admin API) =="
45     RUN_URI="https://${REGION}-
run.googleapis.com/apis/run.googleapis.com/v1/namespaces/${PROJECT}/jobs/${JOB}:run"
46     gcloud scheduler jobs describe "$SCHED" --location="$REGION" --project="$PROJECT"
>/dev/null 2>&1 && \
47     gcloud scheduler jobs delete "$SCHED" --location="$REGION" --project="$PROJECT"
--quiet || true
48     gcloud scheduler jobs create http "$SCHED" --project="$PROJECT" --location="$REGION" \
49         --schedule="$CRON" --time-zone="$TZ" \
50         --uri="$RUN_URI" --http-method=POST \
51         --oauth-service-account-email="$SA"
52
53     cat <<EOF
54
55     ? Scheduled '$SCHED' ($CRON $TZ) ? Cloud Run job '$JOB' ? launch.sh ? ephemeral GPU VM.
56     IAM the launcher SA needs (grant once if missing):
57         roles/compute.instanceAdmin.v1 (create/delete the GPU VM)
58         roles/iam.serviceAccountUser (attach $SA to the VM)
59         roles/run.invoker + roles/cloudscheduler.jobRunner (Scheduler ? Run)
60     Trigger now to test: gcloud run jobs execute $JOB --region=$REGION
--project=$PROJECT
61     Remove: gcloud scheduler jobs delete $SCHED --location=$REGION; gcloud

```

```
run jobs delete $JOB --region=$REGION
62     EOF
63
```

9. user (2026-06-22T09:47:42.765Z)

```
1      # Tiny launcher image for the nightly engine-regression Cloud Run JOB. It is NOT the GPU
workload ?
2      # it only stages the repo + creates the ephemeral GPU VM (which self-deletes). Scale-to-
zero: the job
3      # runs for a few seconds nightly, then exits. Base = google/cloud-sdk (gcloud + git
preinstalled).
4      FROM gcr.io/google.com/cloudsdktool/cloud-sdk:slim
5      WORKDIR /app
6      # The repo is the build context (run `gcloud builds submit` from the repo root, see
register_scheduler.sh).
7      COPY . /app
8      ENTRYPOINT ["bash", "deploy/gcp/nightly_regression/launch.sh"]
9
```

10. user (2026-06-22T09:47:43.574Z)

```
1      {
2      "doc": "Nightly E2E engine reel-count manifest (scripts/nightly_reelcount.py). The
configured+checkpointed pipeline runs on each video on a GCP GPU VM; the reel count + quality
metrics are asserted vs eval_baselines/reelcount_baseline.json. ADD A VIDEO = append one row +
re-freeze: `python scripts/nightly_reelcount.py --update-baseline` (commit the updated
baseline). URIs may be gs:// (recommended, co-located with the VM), gdrive://, or local paths.
labels_uri (optional) is a golden-features JSON (gts key) or a start,end CSV ? present =>
quality metrics (tol_f1/f1@0.5/f1@0.25) are scored too. EDIT the example rows below to your
actual short (~5-6 min) OWNED, minors-free clips before enabling.",
3      "segmenter": "wasb_hybrid",
4      "detector_impl": "native",
5      "config_overrides": {
6      "indexing.default_segmenter": "wasb_hybrid",
7      "indexing.detector_impl": "native"
8      },
9      "machine_usd_per_hr": 0.70,
10     "fx_inr_per_usd": 83.0,
11     "videos": [
12     {
13     "name": "GX030094",
14     "uri": "gs://khelsutra-rally-corpus/nightly/videos/GX030094.mp4",
15     "labels_uri": "gs://khelsutra-rally-
corpus/nightly/labels/GX030094_golden_features.json"
16     },
17     {
18     "name": "Badminton_BXH_2",
19     "uri": "gs://khelsutra-rally-corpus/nightly/videos/Badminton_BXH_2.mp4",
20     "labels_uri": "gs://khelsutra-rally-
corpus/nightly/labels/Badminton_BXH_2_golden_features.json"
21     },
22     {
23     "name": "GX010128",
24     "uri": "gs://khelsutra-rally-corpus/nightly/videos/GX010128.mp4",
25     "labels_uri": "gs://khelsutra-rally-
corpus/nightly/labels/GX010128_golden_features.json"
26     }
27     ]
28     }
29
```

11. assistant (2026-06-22T09:47:46.133Z)

Now let me verify the exact SMTP implementation and search for any other security/correctness issues:

12. user (2026-06-22T09:47:47.916Z)

```
10: ? SMTP host: ``NIGHTLY_SMTP_HOST`` (default ``smtp.gmail.com``) . port  
``NIGHTLY_SMTP_PORT`` (587, STARTTLS)
```

13. user (2026-06-22T09:47:49.070Z)

Help on function `__init__` in module `smtpplib`:

```
__init__(  
    self,  
    host='',  
    port=0,  
    local_hostname=None,  
    timeout=<object object at 0x000001D105BD0360>,  
    source_address=None  
)
```

Initialize a new instance.

If specified, `'host'` is the name of the remote host to which to connect. If specified, `'port'` specifies the port to which to connect. By default, `smtpplib.SMTP_PORT` is used. If a host is specified the connect method is called, and if it returns anything other than a success code an `SMTPConnectError` is raised. If specified, `'local_hostname'` is used as the FQDN of the local host in the HELO/EHLO command. Otherwise, the local hostname is found using `socket.getfqdn()`. The `'source_address'` parameter takes a 2-tuple (host, port) for the socket to bind to as its source address before connecting. If the host is `''` and port is `0`, the OS default behavior will be used.

14. assistant (2026-06-22T09:47:51.244Z)

Let me check Python's `smtpplib` default behavior with `starttls`:

15. user (2026-06-22T09:47:52.308Z)

Help on function `starttls` in module `smtpplib`:

```
starttls(self, *, context=None)  
    Puts the connection to the SMTP server into TLS mode.
```

If there has been no previous EHLO or HELO command this session, this method tries ESMTP EHLO first.

If the server supports TLS, this will encrypt the rest of the SMTP session. If you provide the context parameter, the identity of the SMTP server and client can be checked. This, however, depends on whether the socket module really checks the certificates.

This method may raise the following exceptions:

<code>SMTPHeloError</code>	The server didn't reply properly to the helo greeting.
----------------------------	--

16. assistant (2026-06-22T09:47:54.127Z)

Let me check the Python SSL defaults:

17. user (2026-06-22T09:47:58.311Z)

```

=== Default SSL behavior ===
check_hostname: True
verify_mode: 2
verify_mode == CERT_REQUIRED: True

=== SMTP starttls signature ===
starttls signature: (self, *, context=None)

SMTP.__init__ signature: (self, host='', port=0, local_hostname=None, timeout=<object object at
0x000000192C40B0420>, source_address=None)

=== Testing SMTP behavior ===
When SMTP is created without a context parameter, it uses None
When starttls() is called without a context, Python's behavior depends on version

```

18. assistant (2026-06-22T09:48:00.296Z)

Now let me check Python version and starttls default behavior:

19. user (2026-06-22T09:48:02.155Z)

Python version: 3.13.13 (tags/v3.13.13:01104ce, Apr 7 2026, 19:25:48) [MSC v.1944 64 bit (AMD64)]

```

=== Check Python 3.8+ starttls behavior ===
def starttls(self, *, context=None):
    """Puts the connection to the SMTP server into TLS mode.

    If there has been no previous EHLO or HELO command this session, this
    method tries ESMTP EHLO first.

    If the server supports TLS, this will encrypt the rest of the SMTP
    session. If you provide the context parameter,
    the identity of the SMTP server and client can be checked. This,
    however, depends on whether the socket module really checks the

```

20. user (2026-06-22T09:48:06.252Z)

```

1:     def starttls(self, *, context=None):
2:         """Puts the connection to the SMTP server into TLS mode.
3:
4:         If there has been no previous EHLO or HELO command this session, this
5:         method tries ESMTP EHLO first.
6:
7:         If the server supports TLS, this will encrypt the rest of the SMTP
8:         session. If you provide the context parameter,
9:         the identity of the SMTP server and client can be checked. This,
10:        however, depends on whether the socket module really checks the
11:        certificates.
12:
13:        This method may raise the following exceptions:
14:
15:        SMTPHeloError          The server didn't reply properly to
16:                               the helo greeting.
17:        """
18:        self.ehlo_or_helo_if_needed()
19:        if not self.has_extn("starttls"):
20:            raise SMTPNotSupportedError(
21:                "STARTTLS extension not supported by server.")
22:        (resp, reply) = self.docmd("STARTTLS")
23:        if resp == 220:
24:            if not _have_ssl:
25:                raise RuntimeError("No SSL support included in this Python")
26:            if context is None:

```

```

27:         context = ssl._create_stdlib_context()
28:         self.sock = context.wrap_socket(self.sock,
29:                                         server_hostname=self._host)
30:         self.file = None
31:         # RFC 3207:
32:         # The client MUST discard any knowledge obtained from
33:         # the server, such as the list of SMTP service extensions,
34:         # which was not obtained from the TLS negotiation itself.
35:         self.helo_resp = None
36:         self.ehlo_resp = None
37:         self.esmtp_features = {}
38:         self.does_esmtp = False
39:     else:
40:         # RFC 3207:

```

21. assistant (2026-06-22T09:48:08.753Z)

Excellent. Now let me check what `ssl.\_create\_stdlib\_context()` does:

22. user (2026-06-22T09:48:11.481Z)

```

check_hostname: False
verify_mode: 0
verify_mode == CERT_REQUIRED: False

```

```

create_default_context():
check_hostname: True
verify_mode: 2

```

Are they the same? False

23. assistant (2026-06-22T09:48:13.740Z)

Perfect! This is the actual defect. Now let me verify the exact line numbers and check for other potential issues:

24. user (2026-06-22T09:48:15.198Z)

```

80         if not password:
81             log.warning("no SMTP password (set NIGHTLY_SMTP_PASS or NIGHTLY_SMTP_SECRET) ?
email skipped.")
82             return False
83
84         import smtplib
85         msg = build_message(subject, body_md, from_addr=user, to_addr=to_addr)
86         with smtplib.SMTP(host, port, timeout=60) as s:
87             s.starttls()
88             s.login(user, password)
89             s.send_message(msg)
90         log.info("emailed report to %s", to_addr)
91         return True
92

```

25. assistant (2026-06-22T09:48:17.475Z)

Now let me check for any other correctness/safety issues by examining the deployment flow and credentials handling more carefully:

26. user (2026-06-22T09:48:19.245Z)

```

1: ""Email the nightly regression report ? stdlib SMTP, creds from env or GCP Secret Manager.
2-
3-Used by ``scripts/nightly_reelcount.py --email`` (and reusable by any nightly). Deliberately
4:dependency-light (stdlib ``smtplib`` + ``email.message``) and **graceful**: if no SMTP

```

```

password can
5-be resolved it logs + returns ``False`` (the report is still written/printed ? a broken mailer
must
6-never mask a regression).
--
11- ? SMTP user:  ``NIGHTLY_SMTP_USER`` (default = recipient)
12- ? SMTP pass:  ``NIGHTLY_SMTP_PASS`` ? Secret Manager ``NIGHTLY_SMTP_SECRET``
13:                (``projects/<p>/secrets/<name>/versions/latest``; via the google-cloud-
secret-manager
14:                lib if installed, else the ``gcloud`` CLI). A Gmail **app password** is the
simplest.
15-
16:Set the secret once (owner):  echo -n '<app-password>' | gcloud secrets create nightly-smtp-
pass --data-file=-
17:and on the VM:  export NIGHTLY_SMTP_SECRET=projects/<proj>/secrets/nightly-smtp-
pass/versions/latest
18-                export NIGHTLY_SMTP_USER=<your-gmail>    NIGHTLY_EMAIL_TO=avi.dullu@gmail.com
19-""""
--
42-
43-
44:def resolve_secret(env: Optional[dict] = None) -> Optional[str]:
45:    """Resolve the SMTP password from env or GCP Secret Manager. Returns None if
unresolvable."""
46:    env = env if env is not None else os.environ
47:    direct = env.get("NIGHTLY_SMTP_PASS")
48:    if direct:
49:        return direct
50:    secret_ref = env.get("NIGHTLY_SMTP_SECRET")
51:    if not secret_ref:
52:        return None
53:    # Prefer the client lib; fall back to the gcloud CLI (no extra dependency required).
54:    try:
55:        from google.cloud import secretmanager # type: ignore
56:        client = secretmanager.SecretManagerServiceClient()
57:        return client.access_secret_version(name=secret_ref).payload.data.decode("utf-8")
58:    except Exception: # noqa: BLE001 ? lib absent or auth issue ? try the CLI
59:        pass
60:    try:
61:        import subprocess
62:        name = secret_ref.split("/secrets/", 1)[1].split("/", 1)[0] if "/secrets/" in
secret_ref else secret_ref
63:        out = subprocess.run(["gcloud", "secrets", "versions", "access", "latest", "--
secret", name],
64:                             capture_output=True, text=True, timeout=30)
65:        return out.stdout.strip() if out.returncode == 0 and out.stdout.strip() else None
66:    except Exception as e: # noqa: BLE001
67:        log.warning("secret-manager resolution failed: %s", e)
68:        return None
69-
70-
71-def send_report(subject: str, body_md: str, to: Optional[str] = None,
72:                env: Optional[dict] = None) -> bool:
73:    """Send the report. Returns True if sent, False if skipped (no creds) ? never raises on
missing creds."""
74:    env = env if env is not None else os.environ
75:    to_addr = to or env.get("NIGHTLY_EMAIL_TO") or DEFAULT_TO
76:    host = env.get("NIGHTLY_SMTP_HOST", DEFAULT_HOST)
77:    port = int(env.get("NIGHTLY_SMTP_PORT", DEFAULT_PORT))
78:    user = env.get("NIGHTLY_SMTP_USER") or to_addr
79:    password = resolve_secret(env)
80:    if not password:
81:        log.warning("no SMTP password (set NIGHTLY_SMTP_PASS or NIGHTLY_SMTP_SECRET) ? email
skipped.")
82:        return False

```

```

83-
--
86-     with smtplib.SMTP(host, port, timeout=60) as s:
87-         s.starttls()
88:         s.login(user, password)
89-         s.send_message(msg)
90-     log.info("emailed report to %s", to_addr)

```

27. user (2026-06-22T09:48:20.852Z)

```

40- echo "RUN FAILED ? attempting a failure alert"
41- if [ "$DO_EMAIL" = "true" ] && [ -d ~/rally ]; then
42-     cd ~/rally 2>/dev/null && . .venv/bin/activate 2>/dev/null && \
43:     NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
44-     python - <<PY || true
45-from scripts.nightly_email import send_report
46-send_report("[KhelSutra nightly] ERROR ? engine regression VM run failed ($DATE)",
--
80-UPTIME=$(( $(date +%s) - BOOT_EPOCH ))
81-EMAIL_FLAG=""; [ "$DO_EMAIL" = "true" ] && EMAIL_FLAG="--email --email-to $EMAIL_TO"
82-set +e
83:NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
84- python scripts/nightly_reelcount.py --vm-uptime-seconds "$UPTIME" $EMAIL_FLAG
85-RUN_RC=$?
86-set -e

```

28. assistant (2026-06-22T09:48:23.339Z)

Now let me check for any other potential issues in the architecture. Let me look at the manifest handling, error cases, and broader flow:

29. user (2026-06-22T09:48:25.440Z)

```

54:     try:
58:         except Exception: # noqa: BLE001 ? lib absent or auth issue ? try the CLI
60:             try:
66:                 except Exception as e: # noqa: BLE001
73:                 """Send the report. Returns True if sent, False if skipped (no creds) ? never raises on
missing creds."""

```

30. user (2026-06-22T09:48:26.710Z)

270:# IO shell ? running the configured pipeline per video (needs backend + GPU + video).

31. assistant (2026-06-22T09:48:34.596Z)

Let me check for subprocess usage and environment variable passing in the shell scripts more carefully:

32. user (2026-06-22T09:48:35.197Z)

```

scripts/nightly_email.py:61:         import subprocess
scripts/nightly_email.py:63:         out = subprocess.run(["gcloud", "secrets", "versions",
"access", "latest", "--secret", name],
scripts/nightly_reelcount.py:270:# IO shell ? running the configured pipeline per video (needs
backend + GPU + video).
scripts/nightly_reelcount.py:419:         import subprocess
scripts/nightly_reelcount.py:421:         out = subprocess.run(["git", "-C", REPO_ROOT, "rev-
parse", "--short", "HEAD"],

```

33. assistant (2026-06-22T09:48:37.258Z)

Good! Now let me check for any credential handling in the subprocess calls or environment variable leakage:

34. user (2026-06-22T09:48:38.650Z)

```
60         try:
61             import subprocess
62             name = secret_ref.split("/secrets/", 1)[1].split("/", 1)[0] if "/secrets/" in
secret_ref else secret_ref
63             out = subprocess.run(["gcloud", "secrets", "versions", "access", "latest", "--
secret", name],
64                                 capture_output=True, text=True, timeout=30)
65             return out.stdout.strip() if out.returncode == 0 and out.stdout.strip() else
None
66         except Exception as e: # noqa: BLE001
67             log.warning("secret-manager resolution failed: %s", e)
68             return None
69
```

35. user (2026-06-22T09:48:39.427Z)

```
417
418     def _git_head() -> str:
419         import subprocess
420         try:
421             out = subprocess.run(["git", "-C", REPO_ROOT, "rev-parse", "--short", "HEAD"],
422                                 capture_output=True, text=True, timeout=10)
423             return out.stdout.strip() or "?"
424         except Exception: # noqa: BLE001
425             return "?"
426
```

36. assistant (2026-06-22T09:48:41.930Z)

Now let me check the GCP launcher scripts for any credential leakage in logs or process arguments:

37. user (2026-06-22T09:48:43.235Z)

```
26:SMTP_SECRET=$(md instance/attributes/smtp-secret || true)
43:   NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
83:NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
```

38. assistant (2026-06-22T09:48:45.212Z)

Now let me check the environment variable handling in the startup script for any leakage to logs:

39. user (2026-06-22T09:48:46.342Z)

```
1     #!/usr/bin/env bash
2     # Startup-script for the NIGHTLY engine-regression VM. Runs UNATTENDED on boot, then
DELETES the VM
3     # (cost rule: a running GPU VM is the only real cost). Survives SSH drops (it's a
startup-script, not
4     # an interactive session). Reads its config from instance metadata (set by launch.sh).
5     #
6     # Flow:  Ops Agent ? pull repo @ HEAD from GCS ? build venv + WASB env ? stage
weights ?
7     #         run scripts/nightly_reelcount.py on the gs:// golden clips (emails the report)
?
8     #         upload report to GCS ? SELF-DELETE (even on failure, via the EXIT trap).
```

```

9      #
10     # Metadata keys (launch.sh sets these): repo-tgz-uri, git-sha, gcs-bucket, report-
prefix,
11     # weights-uri, email (true/false). NOT verified in CI ? dry-run via launch.sh on a
real VM first.
12     set -uo pipefail
13     exec > >(tee -a /var/log/nightly-regression.log) 2>&1
14     echo "==== nightly-regression startup $(date -u +%FT%TZ) ====="
15     BOOT_EPOCH=$(date +%s)
16
17     md() { curl -s -H "Metadata-Flavor: Google"
"http://metadata.google.internal/computeMetadata/v1/$1"; }
18     NAME=$(md instance/name)
19     ZONE=$(md instance/zone | awk -F/ '{print $NF}')
20     REPO_TGZ=$(md instance/attributes/repo-tgz-uri || true)
21     GIT_SHA=$(md instance/attributes/git-sha || echo unknown)
22     BUCKET=$(md instance/attributes/gcs-bucket || true)
23     REPORT_PREFIX=$(md instance/attributes/report-prefix || echo nightly/reports)
24     WEIGHTS_URI=$(md instance/attributes/weights-uri || true)
25     DO_EMAIL=$(md instance/attributes/email || echo true)
26     SMTP_SECRET=$(md instance/attributes/smtp-secret || true)
27     SMTP_USER=$(md instance/attributes/smtp-user || true)
28     EMAIL_TO=$(md instance/attributes/email-to || echo avi.dullu@gmail.com)
29     DATE=$(date -u +%F)
30

```

40. assistant (2026-06-22T09:48:49.069Z)

Let me check the logs and output handling, especially around line 13 where logging is set up:

41. user (2026-06-22T09:48:53.182Z)

```

13:exec > >(tee -a /var/log/nightly-regression.log) 2>&1
14:echo "==== nightly-regression startup $(date -u +%FT%TZ) ====="
21:GIT_SHA=$(md instance/attributes/git-sha || echo unknown)
23:REPORT_PREFIX=$(md instance/attributes/report-prefix || echo nightly/reports)
25:DO_EMAIL=$(md instance/attributes/email || echo true)
28:EMAIL_TO=$(md instance/attributes/email-to || echo avi.dullu@gmail.com)
33: echo "==== self-delete $NAME ($ZONE) ====="
35:   echo "WARN: self-delete failed ? DELETE MANUALLY: gcloud compute instances delete $NAME
--zone=$ZONE -q"
40: echo "RUN FAILED ? attempting a failure alert"
48:   "See gs://$BUCKET/$REPORT_PREFIX/$DATE/ and the VM serial
log.\nGIT_SHA=$GIT_SHA", to="$EMAIL_TO")
56: sudo bash add-google-cloud-ops-agent-repo.sh --also-install || echo "WARN: ops-agent
install failed (non-fatal)"
62:echo "$GIT_SHA" > ~/rally/GIT_SHA
70:bash deploy/digitalocean/gpu_setup.sh || echo "WARN: gpu_setup (WASB env) returned nonzero ?
continuing"
87:echo "runner exit code: $RUN_RC (0 pass / 1 regression / 4 stale)"
95:echo "==== nightly-regression done rc=$RUN_RC $(date -u +%FT%TZ) ====="

```

42. assistant (2026-06-22T09:48:56.488Z)

Let me check if credentials are ever logged or echoed, and check all environment variable usages:

43. user (2026-06-22T09:48:57.475Z)

```

NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER"
NIGHTLY_EMAIL_TO="$EMAIL_TO" \
NIGHTLY_SMTP_SECRET="$SMTP_SECRET" NIGHTLY_SMTP_USER="$SMTP_USER" NIGHTLY_EMAIL_TO="$EMAIL_TO" \

```

44. assistant (2026-06-22T09:49:02.395Z)

Good?secrets are passed only as environment variables to the python commands, not echoed. Now let me check if there are any issues with the GCP metadata endpoint or credential leakage:

45. user (2026-06-22T09:49:03.432Z)

```
40     rm -f "$TMP_TGZ"
41
42     # 2) Zone-sweep create the VM (mirrors create_gpu_vm.sh) with run_on_vm.sh + the config
metadata.
43     ERR="$(mktemp)"
44     for Z in $ZONES; do
45         echo "== creating $NAME ($GPU) in $Z =="
46         if gcloud compute instances create "$NAME" --project="$PROJECT" --zone="$Z" \
47             --machine-type="$MACHINE" "${ACCEL[@]}" \
48             --maintenance-policy=TERMINATE --provisioning-model=STANDARD \
49             --image-family=common-cu129-ubuntu-2204-nvidia-580 --image-project=deeplearning-
platform-release \
50             --boot-disk-size=120GB --boot-disk-type=pd-balanced \
51             --service-account="$SA" --scopes=cloud-platform \
52             --metadata-from-file=startup-script="$STARTUP" \
53             --metadata=repo-tgz-uri="$TGZ_URI",git-sha="$SHA",gcs-bucket="$BUCKET",report-
prefix="$REPORT_PREFIX",weights-uri="$WEIGHTS_URI",email="$DO_EMAIL",smtp-
secret="$SMTP_SECRET",smtp-user="$SMTP_USER",email-to="$EMAIL_TO" \
54             2>"$ERR"; then
55             echo "? launched $NAME in $Z ? it will run the regression, email $EMAIL_TO, and
self-delete."
56             echo "    serial log: gcloud compute instances get-serial-port-output $NAME --zone=$Z
--project=$PROJECT"
57             echo "    if it doesn't self-delete: gcloud compute instances delete $NAME --zone=$Z
-q"
58             rm -f "$ERR"; exit 0
59     fi
```

46. assistant (2026-06-22T09:49:06.191Z)

I see on line 53 that `SMTP\_SECRET` and `SMTP\_USER` are passed as GCP instance metadata. Let me verify this is correct and check what happens with instance metadata:

47. user (2026-06-22T09:49:11.215Z)

GCP instance metadata is:

1. Stored in GCP's metadata server (<https://cloud.google.com/compute/docs/storing-retrieving-metadata>)
2. Only accessible from within the VM via <http://metadata.google.internal>
3. Requires the Metadata-Flavor: Google header to access (prevents SSRF via other headers)
4. Default: only readable by the VM's service account (can be restricted further via project policy)

So passing secrets via metadata is a known pattern in GCP and is considered safe IF:

- The metadata server access is restricted (which it is by default)
- The VM has appropriate IAM restrictions (only what it needs)
- The secrets are not passed via gcloud CLI args (which they are on line 53)

This is a potential issue: the gcloud CLI command on line 53 passes the secret as a plain-text command argument, which means it could appear in process listings (`ps aux`) or shell history.

48. assistant (2026-06-22T09:49:13.113Z)

Let me verify this is a real issue by checking GCP documentation and the actual risk:

49. user (2026-06-22T09:49:22.746Z)

GCP Metadata Secret Exposure Analysis:

-
1. launch.sh passes SMTP_SECRET via gcloud CLI args (line 53)
 - Visible in: ps aux during execution
 - Duration: seconds (command completes quickly)
 - Exposure: Medium (anyone with process monitor access)
 2. The secret becomes instance metadata on the target VM
 - Accessible via: metadata.google.internal (requires VM access)
 - But: startup-script runs IMMEDIATELY as root
 - And: VM self-deletes after completion
 3. Better practice: Use gcloud secrets or pass via a secure channel
 - Option A: gcloud secrets create + pass secret ref (current approach)
 - Option B: Pass secret ref via metadata, let VM fetch it (no plaintext pass)

Actual risk level: LOW (but not ideal)

- Secrets are not logged to GCP audit logs as CLI arguments
- Process visibility is limited in Cloud Run environment
- VM is ephemeral and deletes immediately
- BUT: if process monitoring is enabled, secret may be captured

50. assistant (2026-06-22T09:49:24.932Z)

Now let me verify the actual documented finding about SMTP SSL context more precisely:

51. user (2026-06-22T09:49:34.315Z)

=== CRITICAL FINDING: SMTP SSL Context Verification ===

1. Current code (line 86-87):

```
with smtplib.SMTP(host, port, timeout=60) as s:
    s.starttls()
```
2. What happens internally:
 - SMTP.__init__() does NOT take a context parameter
 - starttls(context=None) uses ssl._create_stdlib_context() by default
3. stdlib context (used by SMTP.starttls with no arg):
 - check_hostname: False
 - verify_mode: 0 (0=CERT_NONE, 2=CERT_REQUIRED)
4. create_default_context (recommended for SMTP):
 - check_hostname: True
 - verify_mode: 2 (0=CERT_NONE, 2=CERT_REQUIRED)
5. VULNERABILITY ASSESSMENT:
 - ? STARTTLS is used (encryption is enabled)
 - ? cert verification is DISABLED by default (check_hostname=False)
 - ? An attacker with network access COULD intercept (MITM)
 - BUT Gmail's STARTTLS is strict and will reject unverified certs in practice
6. REAL-WORLD IMPACT:
 - LOCAL NETWORK MITM: Possible but unlikely on GCP/Cloud Run
 - GMAIL SMTP: Gmail servers expect certificate verification
 - PASSWORD LEAK: If MITM succeeds, password sent in plaintext after TLS downgrade
7. CORRECT FIX:

Pass explicit context to starttls():

```
s.starttls(context=ssl.create_default_context())
```

OR use SMTP_SSL instead of SMTP + starttls:

```
with smtplib.SMTP_SSL(host, port, timeout=60, context=ssl.create_default_context()) as s:
```

52. assistant (2026-06-22T09:49:36.698Z)

Perfect. Now let me verify this is a REAL defect in production context by checking Python versions and any other mitigations:

53. user (2026-06-22T09:49:44.109Z)

Python version: `sys.version_info(major=3, minor=13, micro=13, releaselevel='final', serial=0)`

Python version notes:

- Python 3.8+: `ssl.create_stdlib_context()` used by default
- Python 3.10+: Same behavior
- Python 3.13: Same behavior (verified in this env)

Gmail SMTP Server Behavior:

- Port 587 (STARTTLS): Server accepts unencrypted connections
- Supports STARTTLS: Yes
- Enforces certificate verification: Implicitly YES (Gmail is compliant)

PRACTICAL ATTACK SCENARIO:

1. Attacker intercepts connection to `smtp.gmail.com:587`
2. Attacker responds to STARTTLS with self-signed cert
3. Python's `ssl.create_stdlib_context()` does NOT verify the cert
4. Connection 'upgrades' to TLS with attacker's cert
5. Attacker gets plaintext credentials in SMTP LOGIN

CAN THIS HAPPEN IN PRODUCTION?

- GCP Cloud Run: No direct internet access by default (but egress allowed)
- GCP VM: Yes, ephemeral VM has internet access
- BUT: GCP's network is isolated, so local MITM is very unlikely
- Google-owned SMTP: Google's infrastructure is secure

VERDICT: REAL vulnerability, but LOW RISK in Google's cloud environment

However: Best practice would be to explicitly verify certificates

54. assistant (2026-06-22T09:49:46.490Z)

Now let me do a final comprehensive check for other potential defects - checking for any hardcoded credentials, unvalidated inputs, or other security issues:

55. user (2026-06-22T09:49:49.550Z)

```
scripts/nightly_email.py:dependency-light (stdlib ``smtplib`` + ``email.message``) and
**graceful**: if no SMTP password can
scripts/nightly_email.py: lib if installed, else the ``gcloud`` CLI). A Gmail
**app password** is the simplest.
scripts/nightly_email.py:Set the secret once (owner): echo -n '<app-password>' | gcloud secrets
create nightly-smtp-pass --data-file=-
scripts/nightly_email.py:and on the VM: export
NIGHTLY_SMTP_SECRET=projects/<proj>/secrets/nightly-smtp-pass/versions/latest
scripts/nightly_email.py:def resolve_secret(env: Optional[dict] = None) -> Optional[str]:
scripts/nightly_email.py: secret_ref = env.get("NIGHTLY_SMTP_SECRET")
scripts/nightly_email.py: if not secret_ref:
scripts/nightly_email.py: client = secretmanager.SecretManagerServiceClient()
scripts/nightly_email.py: return
client.access_secret_version(name=secret_ref).payload.data.decode("utf-8")
scripts/nightly_email.py: name = secret_ref.split("/secrets/", 1)[1].split("/", 1)[0] if
"/secrets/" in secret_ref else secret_ref
scripts/nightly_email.py: password = resolve_secret(env)
scripts/nightly_email.py: if not password:
scripts/nightly_email.py: log.warning("no SMTP password (set NIGHTLY_SMTP_PASS or
NIGHTLY_SMTP_SECRET) ? email skipped.")
scripts/nightly_email.py: s.login(user, password)
deploy/gcp/nightly_regression/launch.sh:SMTP_SECRET="${NIGHTLY_SMTP_SECRET:-
projects/$PROJECT/secrets/nightly-smtp-pass/versions/latest}"
deploy/gcp/nightly_regression/launch.sh: --metadata=repo-tgz-uri="$TGZ_URI",git-
```

```
sha="$SHA",gcs-bucket="$BUCKET",report-prefix="$REPORT_PREFIX",weights-  
uri="$WEIGHTS_URI",email="$DO_EMAIL",smtp-secret="$SMTP_SECRET",smtp-user="$SMTP_USER",email-  
to="$EMAIL_TO" \  
deploy/gcp/nightly_regression/run_on_vm.sh:SMTP_SECRET=$(md instance/attributes/smtp-secret ||  
true)
```

56. assistant (2026-06-22T09:49:52.246Z)

Let me check for any input validation issues or command injection risks:

57. user (2026-06-22T09:50:02.747Z)

=== PARSING RISK ANALYSIS ===

1. Secret reference parsing (nightly_email.py:62)
Code: name = secret_ref.split('/secrets/', 1)[1].split('/', 1)[0]
Risk: If secret_ref is malformed, could cause IndexError
Impact: Would raise exception (caught by try/except)
Verdict: SAFE (error handling in place)
2. Subprocess calls
 - a) subprocess.run(['gcloud', 'secrets', 'versions', 'access', ...], capture_output=True)
 - Uses list form (not shell=True)
 - No user input in command args
 - Verdict: SAFE
 - b) subprocess.run(['git', '-C', REPO_ROOT, 'rev-parse', 'HEAD'])
 - REPO_ROOT is derived from __file__ (fixed at module load)
 - Verdict: SAFE
3. Email address handling (nightly_email.py)
 - to_addr: from env or default
 - user: from env or defaults to to_addr
 - No regex validation, but used only in:
 - * EmailMessage headers (built-in validation)
 - * SMTP.login() (Gmail validates)
 - * log.info() (safe)Verdict: ACCEPTABLE (Gmail validates)
4. JSON parsing (eval_baselines/reelcount_manifest.json)
 - json.load() is safe (validates JSON structure)
 - Values used in:
 - * config_overrides -> applied to pydantic config
 - * video URIs -> passed to backend.storage.fetch()Verdict: DEPENDS ON DOWNSTREAM (not in scope)

=== OVERALL INPUT VALIDATION: GOOD ===

- No shell injection risks
- No SQL injection (no SQL used)
- JSON parsing is safe
- Credentials handled via env/Secret Manager

58. assistant (2026-06-22T09:50:05.209Z)

Now let me verify there are no other certificate/TLS issues and check for any data leakage in reports or logs:

59. user (2026-06-22T09:50:06.683Z)

(Bash completed with no output)

60. assistant (2026-06-22T09:50:08.980Z)

Good - no credential leakage in nightly_reelcount.py. Let me do a final comprehensive check for any other defects:

61. user (2026-06-22T09:50:23.239Z)

=== COMPREHENSIVE DEFECT ANALYSIS ===

DEFECTS FOUND:

1. SMTP starttls() uses unverified SSL context by default
File: scripts/nightly_email.py:86-87
Severity: low
Real: True
2. SMTP password passed via gcloud CLI arguments (process visibility)
File: deploy/gcp/nightly_regression/launch.sh:53
Severity: low
Real: True

VERIFIED GOOD PRACTICES:

- ? Credentials passed via environment variables (not hardcoded)
- ? Secret Manager integration (graceful fallback to gcloud CLI)
- ? No shell injection (subprocess uses list form)
- ? No SQL injection (no SQL used)
- ? No hardcoded API keys in code
- ? Reports don't leak secrets
- ? Logs don't echo credentials
- ? VM self-deletes immediately
- ? Startup script uses set -uo pipefail
- ? Error handling for missing credentials

FINAL VERDICT:

REAL defect found: SSL certificate verification issue
Severity: LOW (high MITM threshold in Google environment)
Recommended fix: Add context=ssl.create_default_context() to starttls()

62. assistant (2026-06-22T09:50:25.281Z)

Perfect. Now let me create the final structured output with conclusive evidence:

63. user (2026-06-22T09:50:35.115Z)

Structured output provided successfully

64. assistant (2026-06-22T09:50:44.897Z)

Perfect. I have completed the adversarial review of the NIGHTLY REEL-COUNT engine-regression PR. Here's my summary:

Summary

I reviewed the nightly engine-regression system across all specified files:

- `scripts/nightly\_email.py` (SMTP credential handling)
- `scripts/nightly\_reelcount.py` (regression test orchestration)
- `deploy/gcp/nightly\_regression/run\_on\_vm.sh` (VM startup script)
- `deploy/gcp/nightly\_regression/launch.sh` (GCP VM launcher)
- `deploy/gcp/nightly\_regression/register\_scheduler.sh` (scheduler setup)
- `deploy/gcp/nightly\_regression/Dockerfile.launcher` (container image)
- `eval\_baselines/reelcount\_manifest.json` (test manifest)

Finding: CONFIRMED REAL DEFECT

****SSL Certificate Verification Missing in SMTP****

The code at ``scripts/nightly_email.py:86-87`` uses ``smtplib.SMTP().starttls()`` without an explicit SSL context, which defaults to Python's ``ssl._create_stdlib_context()``. This context has certificate verification disabled (``check_hostname=False``, ``verify_mode=CERT_NONE``), unlike the recommended ``ssl.create_default_context()``.

****Why it's real:**** A network MITM attacker could intercept the connection to `smtp.gmail.com:587` and present a self-signed certificate; Python would accept it and the plaintext SMTP password would be compromised.

****Severity: Low**** because:

- Google's infrastructure and network isolation make MITM highly unlikely
- Gmail's SMTP server expects certificate verification in practice
- The risk is conditional on network-level access, not code logic

****Recommended fix:**** Pass ``context=ssl.create_default_context()`` to ``s.starttls()`` for explicit certificate verification.

The rest of the system shows good security practices: credentials via Secret Manager with graceful fallback, no shell injection, no hardcoded secrets, proper error handling, and no credential leakage in logs or reports.