

claude-windows / 386fd981-1117-4be6-8bd0-5ef6dc747c51

Metadata

```
- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\386fd981-1117-4be6-8bd0-5ef6dc747c51.jsonl`
- SHA-256: `a8c8596aaf0c956fe623a1f83310e7bc6bc3e75d18da7b4d9b7744b1f3dd1fb4`
- Source modified: `2026-06-22T09:03:09+00:00`
- Imported at: `2026-07-05T16:48:17+00:00`
- project: `C--Users-avidu-Projects-badminton-highlight-indexer`
- session_id: `386fd981-1117-4be6-8bd0-5ef6dc747c51`
```

Transcript

1. user (2026-06-11T12:52:58.346Z)

In C:\Users\avidu\Projects\badminton-highlight-indexer, the CI workflow .github/workflows/ci.yml uses actions/checkout@v4 and actions/setup-python@v5, which run on Node.js 20. GitHub run logs (June 2026) warn: "Actions will be forced to run with Node.js 24 by default starting June 16th, 2026" and Node 20 is removed from runners September 16th, 2026. Bump to the Node-24-ready major versions (actions/checkout@v5 and actions/setup-python@v6, or whatever the current latest majors are ? verify on the GitHub marketplace), confirm all four CI jobs (Ruff lint, Mypy, import-smoke, full suite with coverage) stay green on a PR, then merge per repo convention (branch off master, one PR, merge only on verified-green statusCheckRollup).

2. assistant (2026-06-11T12:53:04.518Z)

I'll start by reading the current-state memory and the CI workflow file.

3. user (2026-06-11T12:53:06.032Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-469 of 1323 total (59841 tokens, cap 25000). Call Read with offset=470 limit=469 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-reminder>

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-11 (8?9 QUALITY SWEEP IN PROGRESS ? owner approved 8 AIs via
13     selected). **MERGED on green CI: #90** (hygiene: .coveragerc repointed post-#40 +
14     TEST_RUN_SUMMARY ? docs/archives),
15     ***#91** (pydantic .dict()?model_dump, 7 sites, warnings 15?0), ***#92** (? RUFF LANE:
16     ruff.toml E/F/W/B py38 +
17     fix 194 violations + `lint` CI job; **DELETED scripts/generate_khelacharya_flyer.py ?
18     never compiled, line 39 was a
19     literal AI-placeholder sentence**, invisible because nothing swept scripts/), ***#93**
20     (MYPY LANE: mypy.ini lenient
21     green baseline, explicit 7-module quarantine w/ ratchet rule ? main.py+5
22     segmenters+registry; fixed 18 annotation-level
23     errors), ***#94** (? DB FIX surfaced by coverage: `with sqlite3.connect()` = transaction
```

```

scope NOT close ? all 24
19 Database call sites leaked connections; new `_connect()` CM; ResourceWarnings 340?4),
**#95** (COVERAGE FLOOR:
20 full-suite lane runs --cov, **--cov-fail-under=70** vs CI-measured 72% ? ratchet-only),
**#96** (config unknown-key
21 warning: dotted-path [CONFIG WARNING] for typo'd/dropped keys at all nesting levels;
repo config.json verified clean;
22 extra='forbid' still deferred). Suite 349?351 green**. CI = 4 lanes
(smoke/lint/mypy/full+cov). ?? NEXT: **AI-7
23 hybrid-twin collapse** (shared trajectory-hybrid base, fix shot_count drift, dedupe
_wsl_bash, equivalence tests) +
24 **post-AI-7 quality-eval checkpoint** (owner: verify ~0.66-class evals stick; GPU free,
Gemini key ROTATED, live calls
25 OK). Deferred-with-intent this sweep: ruff I/UP rule groups; yolo_hybrid P3/P4 fold-in
(not as follow-up).
26
27 **Checkpoint:** 2026-06-11 (PR #88 ? **#89 MERGED** ? KhelSutra UI-alpha docs bundle on
master `34cf26c`). #88
28 conflicted (branch cut from `feat/async-job-model` pre-squash of #87) ? per
`scratch/PR88_FIX.md`, recreated as
29 `docs/khelsutra-plan-2` off post-#87 master keeping 6 paths (UI_PROPOSAL, HOSTING_PLAN,
UI_ALPHA_EXECUTION_PLAN,
30 mockups/, DECISIONS **ADR-009**, NEXT_STEPS sync); #88 closed + branch deleted. CI
green-gated (statusCheckRollup),
31 squash-merged.
32 - **Dropped from #89, preserved for PR-2 as UNCOMMITTED working-tree edits on master:**
(1)
33 `docs/DEFERRED_HARDENING_PLAN.md` "#16 DONE ? MERGED (PR #87)" status hunk (also saved
as
34 `scratch/PR2_deferred_hardening_status.patch`); (2) `docs/UI_ALPHA_EXECUTION_PLAN.md`
working-agreement rewrite
35 (Cowork designs + leaves `scratch/PRn_HANDOFF.md`; **Claude Code executes all git/PR
mechanics** ? ratified
36 2026-06-11). Don't `checkout --` these away; they commit with PR-2.
37 - ?? Next UI-alpha slice = **PR-2 (upload/download)** per
`docs/UI_ALPHA_EXECUTION_PLAN.md`.
38 - ? **SLATE CLEANED (2026-06-11):** indexer = NO open PRs, master==origin (`34cf26c`),
**all 23 stale local branches
39 pruned** (every one mapped to a MERGED PR #62?#87 before `-D`); only `master` remains.
Collector = NO open PRs, on
40 master. Intentional carryovers (NOT mess): the 2 uncommitted PR-2 doc edits (above),
untracked `scratch/` handoffs,
41 untracked `artifacts/rally_tal_gen0/0.1.0/` (Gen-0 bundle, ADR-008). ?? Collector has
a stray uncommitted
42 `.gitignore` edit (adds `artifacts/` ignore w/ ADR-008 comment) ? mega-session
leftover, owner to commit (micro-PR)
43 or fold into the next collector PR.
44
45 **Checkpoint:** 2026-06-10 (GEMINI BATTERY COMPLETE ? full coverage; owner's billing was
fine, postpay active; ALL
46 failures were burst/capacity artifacts). Merged this stretch: **#80** (chunker re-
encode/downscale, from the chip
47 session ? reviewed+merged), **#81** (rescued the chip branch's orphaned run_signals test
via cherry-pick), **#83**
48 (provider: 6 generate attempts exponential backoff+jitter ? saved chunks live), **#84**
(tracked
49 `eval_baselines/gemini_wrapper_2026-06-10.json`). Suite 336 green.
50 - **FINAL QUALITY TABLE (IoU .5 vs human golden labels):**
51 - **Mahadevpura** (clean-ish, 31 rallies): heuristic 0.657 · Gen-0 0.744 · **two-stage
WASB+Gemini refine 0.83**
52 (P .81 R .84, 32 preds) · **full-video wrapper 0.93** (P 1.0 R .87).
53 - **Kushagra** (multi-court, 32 rallies, FULL 7/7 coverage + audio): heuristic 0.157 ·
Gen-0 0.561 ·
54 **wrapper 0.66** (P .69 R .62, ±1.2s; FPS = background-game pickups + one 53s dead-
time merge ? exactly the

```

```

55         contrast-metric confound).
56     - **STRATEGY (refined):** clean footage commoditized (0.93 cheap); **multi-court drops
the wrapper to 0.66 = the
57     owned model's ship target** (Gen-0 is 0.10 behind at n=6 ? within corpus-growth
reach). Two-stage refine quality
58     (0.83) < full wrapper on short videos; its value = COST on long academy tapes (only
candidate clips uploaded).
59     Wrapper reliability is genuinely flaky (503 waves all session) ? provider-fallback
architecture is load-bearing.
60     - Artifacts: `output/mahadevpura_gemini_refined.csv`,
`output/MahadevpuraSingles_highlights.mp4` (demo reel),
61     Kushagra preds under the ORIGINAL corpus MD5 now. Total Gemini spend today: ~$1-2.
62     - PENDING/owner: rotate the chat-pasted key; ship-gate F1@tIoU decision (multi-court
target now has a number: 0.66);
63     C13 interviews + camera pilot; repo rename.
64
65     **Checkpoint:** 2026-06-10 (GEMINI BASELINE LANDED ? owner's "gradual ramp-up" theory
CORRECT). New key in gitignored
66     `.env` (rotate later ? pasted in chat). The "credits depleted" 429 was a BURST-THROTTLE
artifact: text ping + serial
67     upload worked ? **PR #79 merged**: gemini segmenter now honors `indexing.ai_max_workers`
(was hardcoded 5; config ships
68     1=serial ? at 5 workers the same video 429'd, at 1 it worked).
69     - ? **WRAPPER BASELINE (gemini-flash-latest, ~$0.2-0.5 total): Mahadevpura F1 0.93** (P
1.00 R 0.87, mIoU 0.80, ±0.7s;
70     3 of 4 misses = short rallies Gemini FOUND but our min_segment_duration=3.0 rejected)
vs Gen-0 learned 0.744 vs
71     heuristic 0.657. **Clean footage = commoditized.**
72     - ?? **MOAT TEST (Kushagra multi-court, windowed n=12 GT): wrapper F1 0.56** (P .54 R
.58 mIoU .69) ? Gen-0 learned
73     0.561, both >> heuristic 0.157. **The multi-court confound hurts Gemini too ? the
wedge survives on the target
74     footage class.** Caveats: only chunks 2+4 answered (5/7 lost to GENUINE gemini-flash-
latest capacity 503s ? retry
75     off-peak for full coverage); visual-only (audio stripped in downscale); 1280p re-
encode (4K chunks were 1.3GB >
76     500MB upload cap ? the audit's silent-oversize-skip fired live; fix = chip
task_44e72c41 ? see PR #80 checkpoint
77     below). Kushagra preds keyed under downscaled-file MD5 (`output/kushagra_1280.mp4`).
78     - ? **DEMO ARTIFACT: `output/MahadevpuraSingles_highlights.mp4`** ? real 27-rally reel
compiled by VideoCompiler
79     (the once-broken module) from Gemini segments. The academy-pitch object exists.
80     - **STRATEGY READ:** product path = Gemini for clean footage TODAY (cheap, F1 0.93) +
owned model's job = multi-court
81     (0.56 is the bar to beat there, not 0.93); 503 flakiness = real reliability cost of
the wrapper (provider-fallback
82     story strengthens). PENDING: gemini_refine two-stage (503-gated, run off-peak), full
Kushagra re-run, PR #80 merge.
83
84     **Checkpoint:** 2026-06-10 (4K gemini oversize fix ? **MERGED: #80 + #81 + #82; master
`f01387e`, CI rollup SUCCESS**).
85     Root cause of the live "0 segments / success" on KushagraYashAviFirstVideo.MP4 (4K):
gemini segmenter stream-copies 90s
86     chunks ? all 7 at 832?1299MB > provider `MAX_UPLOAD_MB=500` ? every chunk silently
skipped. Fix (3 layers): (1) chunks
87     RE-ENCODED + downscaled to new declared knob `indexing.ai_chunk_scale_width` (default
1280 = gemini_refine's
88     `--clip-scale-width`; 0 = legacy stream-copy; also fixes keyframe drift); (2) skip is
LOUD ? provider
89     `metrics["oversize_skipped"]` + segmenter aggregation (chunked + single-file) ?
`logger.error` + NEW thread-safe
90     `backend/reporting/run_signals` channel (incr/pop by video_id; popped unconditionally in
`emit_indexer_report`);
91     (3) report flags it ? `ai_handoff.details.oversize_clips_skipped` + spec.yaml ERROR rule
`oversize_clip_skipped`.

```

```

92 Suite 325? **336 green** locally (CI runs fewer: reporting tests importorskip obsreport,
absent there BY DESIGN ?
93 `tests/test_run_signals.py` deliberately has NO importorskip = the CI-visible half of
the channel). **Merge dance
94 (process note):** owner merged #80 mid-session (auto-deleting the remote branch) WHILE
the follow-up test commit was
95 being pushed ? push re-created the branch orphaned; owner rescued it as #81, 39s before
my cherry-picked #82 ?
96 **82's squash `f01387e` = EMPTY commit** (harmless dup in history). Lesson: re-check
`gh pr view` state before pushing
97 more commits to a PR branch ? owner + sibling sessions work concurrently (this
checkpoint file too: another session
98 prepended the GEMINI-BASELINE checkpoint mid-edit). ? **OWNER/SIBLING WIP IN TREE
(uncommitted, do NOT clobber):**
99 `backend/providers/gemini.py` ? generate retries 3?6 w/ exponential backoff + jitter
(~3?96s) for the 503 capacity
100 waves. ?? commit/PR the 503-backoff WIP . machine-side 4K re-run (chunks should be ~tens
of MB) . task-17 battery
101 still pending Gemini billing top-up.
102
103 **Checkpoint:** 2026-06-10 (M0 GREENLIT SEQUENCE EXECUTED ? session wrap). Owner
ratified ADR-008 ("split driven by
104 productionization, not isolation") + greenlit M0.4+M0.1. **All 5 steps shipped, every PR
merged on verified-green CI:**
105 - **73 ADR-008** (DECISIONS.md): training in-repo behind an executable wall; artifact-
bundle boundary; ship-gate
106 product-side; pre-committed split trigger ? `sports-temporal-models` at Gen-2/second-
consumer.
107 - **74 canonical GT loader** `backend/eval/gt_loader.py::load_gt_intervals` ? closes
the two-loader fork (accepts BOTH
108 `start,end` and `start_time,end_time`); legacy loaders delegate; round-trip contract
test; LOVO 0.626 re-verified.
109 - **75+76 featurizer promotion + import wall** : `backend/features/rally_seq.py`
(FEATURE_SCHEMA v1, fps semantics
110 explicit) imported by train AND serve; `training/` skeleton (5 wall rules,
requirements-train.txt); smoke tests:
111 runtime no-training-in-backend.main, static sweep, featurizer purity (torch/cv2
blocked; importorskip numpy on the
112 minimal CI lane ? #76 hotfixed a red lane after #75 merged early; merges now gated on
verified SUCCESS rollout).
113 - **collector #19 ? n=6 OWNED CORPUS IN THE SoR** : all 6 golden matches via `import-
local --normalize` (262 rallies,
114 0 dropped, Proprietary, explicit fps); commercial export now 6 real matches + demo
entry.
115 - **77 Gen-0 harness** `python -m training.gen0.harness --manifest ? --floor
eval_baselines/heuristic_lovo_n6.json
116 --version 0.1.0` : train?honest LOVO?ship-gate?artifact bundle (gitignored weights.json
+ manifest.json w/ schema pin,
117 calibration, lineage gt_hashes, attestations incl. C3, gate verdict). **First real
bundle built: LOVO 0.626, floor
118 PASSED (0.480), paired sign test 5/6 wins p=0.109 = honestly NOT significant,
ship=False** (blockers: target
119 undefined, C3, significance). Committed floor baseline
`eval_baselines/heuristic_lovo_n6.json`. Suite **325 green**.
120 - ? **Gemini baseline STILL BLOCKED after owner provided a NEW key (2026-06-10 later):**
key written to gitignored
121 `.env`, validated via `models.list` (55 models incl. `gemini-flash-latest`; **config
migrated off deprecating
122 gemini-2.5-flash ? `gemini-flash-latest` alias, PR #78 merged** ; gemini_refine already
defaulted to it). But inference
123 = 429 RESOURCE_EXHAUSTED "prepayment credits depleted" ? **the AI-Studio PROJECT
behind the key has no credits** (key
124 auth ? billing). Owner: add credits at https://ai.studio/projects (+ consider rotating
the chat-pasted key). Planned
125 battery on unblock (in task 17): wrapper Mahadevpura ? score; wrapper Kushagra (multi-

```

```

court MOAT test); two-stage
126     gemini_refine; compile a real highlight reel (academy-demo artifact). Pipeline handled
both failed runs gracefully
127     (validators pass, reports emitted, artifacts cleaned); shared-genai-client thread-
safety finding observed live.
128     - **C13 field plan (owner executing):** recruit interviewer; camera idea RATIFIED w/ 3
constraints (consent-at-capture/
129     adults-only for training; demo via Gemini-or-human-polish, NOT gated on owned model;
interviews?cameras sequencing).
130     - ?? NEXT: owner tops up Gemini billing (wrapper baseline) · owner sets the ship-gate
F1@tIoU target · grow corpus
131     (camera pilot videos ? import-local ? re-run harness) · repo rename · M0.1 event-index
contract in collector.
132
133     **Checkpoint:** 2026-06-10 (Roora PB+TT pilot review ? REPORT DELIVERED to owner; owner
emails Roora himself) ·
134     ? Owner cross-checked the rally tables in VLC ? **CONFIRMED: every boundary off ?0.5s
(up to ~2s)** vs the ±0.3s bar
135     (Brief §4). **Root cause (my forensics): all 24 boundaries sit EXACTLY on the 30-frame
sampling grid** (~0.5s at
136     59.94fps) ? first/last box never on the true contact/dead frame ? structurally cannot
meet ±0.3s. Also decoded their
137     bogus `start/end time` attrs = **frame+900 (decimal minutes @ assumed 15fps, ~4× off
wall-clock)**. Report findings:
138     boundary accuracy (CRITICAL), skip-rallies-in-progress-at-file-start/end convention (PB
rally#1 starts at frame 0 ?
139     shouldn't have been labeled), dead-time+warmup reminders for Phase-B, schema cut to
exactly
140     `rally_number,shot_count,ending_reason,last_shot_outcome,note` (drop their time attrs +
scores), field quality
141     (PB outcome='0' off-vocab; TT 5/6 'other' + missing required notes; TT#5 intra-rally
other|winner). Per owner:
142     NO redelivery demand in the report (he words the email himself; thread = Jamshad@roora
on avidullu@live.com Outlook).
143     **Artifacts:** PDF `Annotation Setup\6 June 2026\Feedback\Roora Pilot Review -
Pickleball and Table Tennis
144     (2026-06-10).pdf` + md record at `Annotation Setup\Roora Pilot Review ? Pickleball &
Table Tennis (2026-06-10).md`
145     (generator: `%TEMP%\roora\make_review_pdf.py`). ? Local guideline docs are **DRAFT v2**
? owner mentioned "v7";
146     no v7 found on disk (possibly in the email thread; report cites v2). Spawned chip
`task_c55f0be6`: collector
147     `rally_cvat.py` ATTR_MAP drops `shot_count` (present in Roora XML, empty in our CSVs) ?
OUR bug, kept out of the
148     vendor report. ?? owner sends email; Roora's badminton long-video + redelivery decisions
pending owner.
149
150     **Checkpoint:** 2026-06-10 (REPO-TOPOLGY decision panel + C13 plan ? AWAITING OWNER
RATIFICATION). Owner leaning
151     M0.4+M0.1 greenlight; asked: model in-repo vs separate repo consumed as a segmenter
(owner prefers isolation). Ran a
152     4-agent judge panel (`wf_471a9cdd-d1d`; full output `?\tasks\wnksv545g.output`).
**Verdict (my rec): "artifact boundary
153     NOW, repo split at Gen-2"*** ? M0.4 lands as a top-level `training/` tree in the indexer
w/ executable isolation
154     (import-smoke asserts backend.main loads no training module; own requirements-train.txt
+ 3rd CI lane), featurizer
155     promoted to `backend/features/rally_seq.py` (single source, train==serve), artifact
bundle = gitignored weights
156     (private Releases) + committed manifest (feature-schema ver, normalization, calibration,
corpus ref+gt_hash,
157     consent/license attestation incl. C3), ship-gate stays product-side (regression.py
extension); **pre-committed split
158     trigger** to a named repo (`sports-temporal-models`) when Gen-2/cloud-training or a 2nd
artifact consumer arrives.

```

```

159     Decisive evidence: Gen-0's input = the indexer's OWN featurization (mid-pipeline seam ?
WASB/Gemini video-boundary);
160     indexer not pip-installable ? split today = copy-paste = the observed drift class;
obsreport "precedent" never actually
161     exercised (pin is a commented line, repo unpublished, CI skips). **FIRST contract to
freeze: golden-label/corpus
162     contract** ? indexer carries TWO incompatible GT loaders (annotations.py `start,end` vs
calibrate_wasb `start_time,
163     end_time` ? fed to 7 modules) vs collector export `start,end`; consolidate + round-trip
test BEFORE any training epoch.
164     Panel extras: n=6 corpus's durable home = loose `Annotation Setup\` folder (neither
repo!) ? M0.1 first deliverable =
165     hash-addressed in collector; training fixtures synthetic-only until consent ledger
exists; run the Gemini-wrapper
166     baseline before/with M0.4. **C13 field plan:** owner recruiting an interviewer; my rec =
bake in the camera idea
167     (paid recordings = consented owned multi-court corpus + demo), demo served via Gemini
path or human-polish (owned model
168     honest 0.626 ? >0.8 yet; the comparison doubles as the wrapper baseline). ?? awaiting
owner ratification of topology +
169     M0.4/M0.1 scope; then: ADR + GT-loader consolidation + featurizer promotion + harness
scaffold.
170
171     **Checkpoint:** 2026-06-10 (CORPUS RECLASS EXECUTED + NEXT_STEPS refreshed ? session
wrap). Owner approved + I ran the
172     corpus correction: `reclassify-licenses --apply` (23 ShuttleSet broadcast rows ?
annotation_license=MIT, video=Unknown,
173     is_commercial=False; **zero annotations deleted**, 1,917 rallies intact) + `curate
export` (commercial manifest now
174     owned-data-only: 1 video/2 segments; stale ShuttleSet CSVs removed from the gitignored
eval_export; eval-only re-export
175     available via `--include-staging --include-noncommercial`). **Merged: collector #17**
(DB + .gitignore backup rule);
176     pre-change backup at `data/sports_metadata.backup-2026-06-10.db` (gitignored). Quality
signals UNAFFECTED (n=6 corpus
177     never flows through the collector DB; learned 0.626 re-verified earlier). **Indexer #72
merged:** NEXT_STEPS.md
178     refreshed (sweep banner, 0.626 vs 0.480, M0.3 partially-done with corrected
distill_local pointer, NEW prioritized PMF-
179     validation item, durable artifact paths). **Owner is now reading NEXT_STEPS to chalk out
next steps** ? open owner
180     decisions stand: greenlight M0.x (rec M0.4+M0.1), confirm collector=SoR, ship-gate
F1@tIoU, repo rename, C13 interviews
181     / Gemini-wrapper baseline. Both repos clean on master, 0 open PRs.
182
183     **Checkpoint:** 2026-06-10 (HIGH-RISK REMEDIATION COMPLETE ? owner approved "everything
A?K, open+merge PRs after CI
184     green, re-run quality evals, GPU available"). **ALL 11 PRs MERGED + CI green; both repos
clean on master.**
185     - **Indexer (badminton-highlight-indexer):** #63 compiler-syntax fix+CI (prior), **#64**
detector keying by video_id,
186     **#65** re-ingest destroy-after-confirm (add_video UPSERT + watermark/prune), **#66**
nested-config wasb typed,
187     **#67** eval gate + metric footguns (iou=0, midrank AUC, fps-normalized speed,
baseline?`eval_baselines/` exit-3),
188     **#68** API path validation, **#69** reliability (timeouts 1800s, SQLite
WAL+busy_timeout, chunk-step guard, CORS),
189     **#70** test coverage (TestClient endpoints + compiler + WasbRunner), **#71** docs
truth-sync. Suite 256?308 green**.
190     - **Collector (sports-data-collector):** **#14** licensing gate (video vs annotation
license; `reclassify-licenses`
191     found 23 conflated ShuttleSet rows) + **added collector CI**, **#15** import-local
golden-label guard (+`--force`,
192     `--normalize`, duplicate-rally guard), **#16** validate-delivery fails on
empty/corrupt/zip CVAT. Suite 124?136 green**.

```

```

193 - ? **QUALITY RE-VALIDATION (owner asked): metric fixes STRENGTHENED the conclusion.**
After midrank-AUC +
194     fps-normalized speed + occlusion-gap guard, **learned LOVO 0.583 ? 0.626** (per-video
AUC unchanged 0.83?0.92 =
195     never tie-inflated; fps-fix mainly helped the 30fps short fold 0.30?0.56). Heuristic
LOVO **0.480** unchanged (scored
196     at IoU 0.5; my iou>0 guard is a no-op there). **Learned beats heuristic by +0.146**
(was +0.10).
197 - **Eval artifacts RESCUED from `%TEMP%` ? durable `C:\Users\avidu\Projects\Annotation
Setup\Collect\Trajectories\`**
198     (+ `\roora\`); manifest `output/human_lovo_manifest.json` repointed there + reverified
0.626.
199 - **Deferred WITH INTENT (in BACKLOG): repo-rename; Gen-2 cost benchmark; ship-gate
F1@tIoU target; packaging
200     (pyproject); app-factory/lifespan + lazy torch-cv2 imports (with async slice #16); WSL
temp/frame/cache + Gemini
201     remote-file GC (retention policy); shlex-quote WSL config values (breaks `~` expansion
? non-trivial); deeper
202     validator tests + 3-hybrid contract test (partial coverage shipped).**
203 - ? **Process note:** during PR-F I briefly `git checkout --`'d the owner's uncommitted
`data/sports_metadata.db`
204     change (200704b), then RECOVERED it from dangling stash `d3c2e77` and restored it to
the working tree (no loss).
205     Going forward stash-and-pop that DB for collector PRs; never `checkout --` it. The
corpus-DB reclassify (`curate
206     reclassify-licenses --apply` + `curate export`) is left for the OWNER to run
deliberately (binary SoR data).
207 - Plan doc (reference): `scratch/HIGH_RISK_REMEDIATION_PLAN.md` (untracked). Audit JSON:
`?\tasks\wf38bovts.output`
208     (batch 1) + `?\tasks\wa4d2980l.output` (full 8 auditors).
209
210     **Checkpoint:** 2026-06-10 (compiler SyntaxError FIX ? **PR #63 MERGED**, squash `master
13d4573`; owner approved merge-on-green) .
211     ? **master UNBROKEN again ? fixed the ? SyntaxError from 7fbbeb0** (= the audit's fix
chip `task_2b7a19d4` work):
212     restored the f-string quotes on the 7 mangled lines of
`backend/pipeline/stitcher/compiler.py`
213     (100/104/156/160/181/191/213; message texts verified against the pre-conversion file at
`7fbbeb0^`; line 191 =
214     plain string, no placeholders) + fixed line 168 `f"\nSUMMARY]" ? `f"\n[SUMMARY]`. **NEW
regression guard
215     `tests/test_import_smoke.py`**: (a) py_compile sweep over every file in backend/
(imports nothing ? runs on
216     cv2-free machines), (b) subprocess `import backend.main` with missing cv2/torch/numpy
stubbed. Negative check:
217     the broken compiler.py fails the sweep (`SyntaxError: invalid decimal literal`, line
100).
218     ? **CI NOW EXISTS ? the audit's no-CI finding is CLOSED** (`github/workflows/ci.yml`,
runs on every PR + master
219     push): job 1 **import-smoke** = smoke tests on a GPU-stack-free runner (minimal dep set
verified empirically in a
220     fresh venv: pytest/fastapi/uvicorn/pydantic/python-dotenv/**google-genai** ? needed at
import time by
221     backend.providers); job 2 **full-suite** = requirements.txt with **CPU-only torch
wheels**
222     (`--index-url download.pytorch.org/whl/cpu`) + libgl for opencv ? `python
run_all_tests.py` (obsreport absent ?
223     reporting tests skip by design). **Both jobs GREEN on the first run** (27s / 1m29s).
Verified post-merge:
224     `import backend.main` OK on master, suite 256 green (254+2).
225     ?? Audit's **coverage-inversion** finding (wasb_hybrid/validators untested) remains open
? CI only guards the
226     import/syntax class + whatever the suite covers.
227
228     **Checkpoint:** 2026-06-10 (DEEP AUDIT session ? ultracode 18-agent audit

```

```

`wf_2a64bc88-2a8` over indexer+collector:
229     engineering/testing/reliability/docs-consistency/PMF; 3 of 8 auditors
completed+verified, 5 rate-limited).
230     - ? **CRITICAL ? master is BROKEN since 7fbbbe0 (2026-06-07 print?logger sweep):**
231     `backend/pipeline/stitcher/compiler.py` has **7 unquoted f-string lines
(100/104/156/160/181/191/213)** ? SyntaxError ?
232     `import backend.main` fails ? FastAPI server + `/api/compile` dead. Triple-verified +
reproduced first-hand
233     (`py_compile` exit 1). The **254-green suite never noticed** (zero tests import
backend.main / compiler / wasb_hybrid /
234     motion / any concrete validator; **no CI exists at all**). Spawned fix chip
`task_2b7a19d4` (fix + import/compileall
235     smoke test). Coverage is INVERTED vs production risk (tracknet twin 8 tests vs live
wasb_hybrid 0; TrackNetRunner 21
236     vs WasbRunner 1). Suite itself healthy where it exists (254 green / 19.9s / faithful
mocks / low flake risk); "fully
237     mocked runs anywhere" is false on cv2-free machines (bare `import cv2` at collection
time, only obsreport importorskip).
238     - **Doc drift (verified, one-directional ? newest decisions not back-propagated):** (1)
licensing "Apache-2.0 only" still
239     in CLAUDE.md:33 / PLATFORM:269,301 / COMMERCIALIZATION C14 vs ratified
**MIT/Apache/BSD** (a literal CLAUDE.md reader
240     rejects the plan's own MIT ActionFormer/ASFormer heads); (2) **Gemma-4 still
"approved" in those same docs vs
241     AMBER/bench-only in the plan; (3) **M0.3 purge pointer WRONG** ?
`training.label_candidates` is a pure GT-agnostic fn;
242     the live Gemini-silver training violation = `backend/eval/distill_local.py` (+
calibrate_local thresholds) + gitignored
243     `output/local_rally_model.json`; (4) **C3 escape hatch broken** ? "train-own weights
(Phase-4 roadmap)" cites a phase
244     that contains no such plan and ADR-002 says corpus labels are temporal-only (CAN'T
train a detector) ? own-weights is
245     an unscoped new workstream (coordinate labels or detector swap); (5)
OWNED_MODEL_IMPLEMENTATION_PLAN + NEXT_STEPS are
246     linked from NO entry point (CLAUDE.md "read first" / docs/README index) and NEXT_STEPS
is branch-only; (6) README
247     "yolo_hybrid by default" vs config.json `gemini` vs CODE_MAP "wasb = LIVE default";
(7) ROADMAP still says audio
248     "gated on golden labels" (trigger now satisfied!) vs plan "parked, not in reserve";
(8) Roora multisport CSVs live ONLY
249     in `%TEMP%` (OS-clearable) + n=6 manifest gitignored ? headline
TT-0.909/pickleball-0.308 evidence is one cleanup from
250     loss; (9) TT-on-WASB idea silently walks into the C3 ship-gate ("multiplies per
sport"); (10) ship-gate F1@tIoU still
251     undefined while 0.44/0.54/0.480 floors circulate; repo-rename "ASAP" tracked in no
tracker (BACKLOG lacks it).
252     - **PMF (web-grounded, sources in audit output):** pricing/cost benchmarks VERIFIED
(SwingVision $179.99/yr ~20k paying
253     subs rev $2.75M +12% YoY; PB Vision $99/yr + **$8/hr-of-video partner API**; Gemini
COGS ~$0.30?0.60/hr ? fat margin);
254     but **"India badminton whitespace" FALSIFIED ? Bangalore is the most contested
badminton-AI city** : Machaxi ($1.5M
255     Rainmatter+Prakash Padukone), VISIST.AI (70+ Bengaluru academies, Padukone School),
Game Theory (Gopichand-backed smart
256     courts w/ auto-highlights), Stupa ($3.3M, BWF-approved, TT?badminton) ? NONE in our
2026-06-07 competitive table.
257     Surviving wedge (narrow, real): **software-only rally-precise indexing of EXISTING
messy multi-court BYO footage**
258     (market solves it only with per-court hardware) sold per-hour as engine/API, possibly
THROUGH those players
259     (PB Vision×PodPlay template). Fastest falsification: C13 interviews (still undone) +
**Gemini-wrapper baseline vs the
260     n=6 LOVO corpus** (if wrapper ? learned 0.583, owned-model is premature) + 1 paid
academy pilot.
261     - **AUDIT NOW COMPLETE (8/8 auditors, resumed `wa4d29801` after Max-20x refresh; 40

```



```

agents).** New high/critical findings
262     beyond the first batch:
263     - ? **CRITICAL (collector) ? LICENSING GUARDRAIL VIOLATION:** `src/cli/ingest.py:12`
defaults `--license MIT`; the
264     ShuttleSet **annotation** license is stamped onto **YouTube broadcast video** (e.g.
shuttleset_21 = a YONEX Thailand
265     Open broadcast, `is_commercial=true, license_type=MIT`). **24 videos / 1915 segments
currently exported as
266     "commercially-safe" on unsound grounds** ? directly violates the "training data must
be commercial-clean" guardrail.
267     - **HIGH (indexer, all verified):** (a) **nested config sections silently drop unknown
keys** ? only `AppConfig` has
268     `extra='allow'`, so the ENTIRE `indexing.wasb` block
(weights/conda_env/velocity_threshold) + `yolo_prefetch_workers`
269     + `debug_duration` are unconfigurable via config.json; meanwhile
`yolo_parallel_chunks/yolo_chunk_workers` are read
270     nowhere. (b) **detector artifacts keyed by file BASENAME not video_id** ? two
`match.mp4` files reuse stale frames ?
271     a trajectory + candidate_segments computed from the WRONG video (silent training-
data poisoning). (c) **re-ingest is
272     destroy-before-confirm** ? `add_video('processed'/'failed')` INSERT-OR-REPLACE
cascade-deletes prior good segments
273     BEFORE the new run; a failed/empty re-run leaves zero. (d) **stitching padding
config not wired** ? API reels always
274     end_padding=1.0 despite config 0.5. (e) **`/api/process` arbitrary server-file read
+ `/api/compile` path-traversal/
275     absolute write** via output_filename. (f) **run_regression "CI gate" not in CI +
silently passes** on missing baseline;
276     baseline omits iou/detector/GT-hash. (g) **metrics bug**: `match_intervals` at
iou_threshold=0.0 matches
277     non-overlapping intervals (P=R=F1=1.0 for garbage); **rally_seq_proto AUC has no tie
handling** (order-dependent;
278     all-zero dead-time rows guarantee ties) + speed features fps-inconsistent within one
vector.
279     - **HIGH (collector):** silent DELETE-then-INSERT replace of golden labels (no
warn/backup, filename-derived video_id);
280     validate-delivery **PASSES on a zero-rally CVAT export**; not yet SoR (no
provenance/maturity/labeled_window cols,
281     import-local hardcodes CURATED); `int(float())` rally_number crashes mid-import
leaving a committed video row w/
282     rolled-back annotations (normalize_annotations guard is opt-in, NOT wired into
import-local).
283     - **MEDIUM batch:** WSL `timeout_sec=0` = hang-forever (both runners + provider
PROCESSING poll); `uvicorn backend.main:app`
284     leaves globals None (init only in main()); SQLite no WAL/no busy_timeout + swallow-
on-lock; chunk_overlap?duration =
285     infinite loop; WSL temp/frame/cache + Gemini remote-file leaks; CORS
`**+credentials; sport-profile relevance config
286     dead under typed config (bites first non-badminton sport ? the current branch!).
287     - **VERDICT NUANCE:** verifier DOWNGRADED 2 "high"?not-real (Result-contract gap + DB-
swallow are DOCUMENTED CODE_MAP
288     gotchas/deliberate-incremental, medium at most; DB "no busy_timeout" claim factually
wrong ? sqlite3 default
289     timeout=5s). Eng-pipeline/eval/collector strengths confirmed: registry seams,
DetectorRunner ABC, WASB crash-cache,
290     honest fit-on-self labeling, leakage-free LOVO, frame-true CVAT conversion all
genuinely good.
291     - ? **SPOOF-CHECK PASSED ? quality iterations are REAL (re-ran the evals first-hand on
this box, pure numpy/scipy):**
292     all 6 videos + 6 human label CSVs + 6 trajectories exist; traj row-counts match docs
exactly. **`rally_seq_proto` LOVO
293     reproduced: mean 0.583** (AUC 0.877/0.851/0.916/0.832/0.730/0.902 vs doc
0.879/0.853/0.915/0.833/0.726/0.902).
294     **`calibrate_local` heuristic LOVO reproduced: 0.480 ± 0.242** (per-fold Adarsh
0.751/Kushagra 0.157/doubles 0.737/

```

295 gBaaddy 0.309/testlarge 0.267/mahadevpura 0.657 ? exact). **Contrast metric recomputed independently from raw CSVs:**

296 3.7/1.4/3.7/1.9/1.4/2.2x vs doc 3.9/1.4/3.6/1.9/1.4/2.2. Opus 4.8 did NOT spoof. ? one nit: gBaaddy CSV has 48 rallies

297 (doc says 47). Caveat: trajectories were WASB-GPU-generated (can't regenerate here) but sizes/frame-counts/visibility all sane.

298 - **HIGH-RISK REMEDIATION PLAN written for owner approval:**

299 PR-1 compiler fix already DONE via #63). Awaiting owner scope + push/PR decision before execution. Full audit JSON:

300 `%LOCALAPPDATA%\Temp\claude\...\tasks\wa4d2980l.output` (8 auditors) + `wf38bovts.output` (first batch).

301 Full audit JSON: `%LOCALAPPDATA%\Temp\claude\C--Users-avidu-Projects-badminton-highlight-indexer\45ea3a91-654d-4972-9cd3-65f35db54caf\tasks\wf38bovts.output`.

302

303 **Checkpoint:** 2026-06-09 (Rooru MULTISPORT recon ? for FRESH next session) . ??

304 **Pickleball + Tabletennis golden labels

304 arrived; INGESTION = SOLVED, but the "merged model" experiment is BLOCKED on per-sport trajectory detectors (validates the

305 roadmap blocker).** Owner left a recon-only ask (deliberately did NOT drain context / spin heavy jobs ? owner wants me fresh).

306 - **Deliverables:** `Annotation Setup/6 June 2026/{pickleball,tabletennis}/\*.zip` (3 CVAT export FORMATS each: COCO

307 `instances\_default.json`, **CVAT-XML `annotations.xml`** ? use this, Pascal-VOC).

308 Videos: `~/6 June 2026/Annotation Videos/`

308 Video 2 - Pickleball.mp4` (4K, 59.94fps, 97s, 5737fr) + `Video 3 - Tabletennis.mp4` (4K, 59.94fps, 58s, 3490fr). Short clips.

309 - **What Rooru annotated:** label class "rally" as **per-frame BOUNDING BOXES on sampled frames (~every 30fr)** ? box geometry

310 is just a CARRIER; the boxes carry the FULL rally schema as CVAT **attributes** (rally_number, start/end time, shot_count,

311 score_before/after(A-B), ending_reason, last_shot_outcome [+note for TT]). 456/513 attr instances. Boxed frames cluster into

312 rallies (runs of ~30-spaced frames = a rally; big gaps = dead time).

313 - **INGESTION = YES, cleanly:** the collector's `src/parsers/cvat/rally\_cvat.py` (`curate convert-cvat`) is PURPOSE-BUILT for

314 exactly this ? ignores box geometry, **recomputes rally interval from the frame range** (`sec=frame/fps` at true fps, ignoring

315 vendor start/end times which are unreliable), groups boxes by rally_number, maps attrs?canonical. Minor: verify the attr-name

316 mapping handles spaces/"(A-B)" (`ATTR\_MAP`). Collector already has pickleball + tabletennis schemas.

317 - ? **"MERGED LEARNT MODEL" across sports is BLOCKED:** `rally\_seq\_proto` features come from a WASB **shuttle** trajectory =

318 BADMINTON-ONLY; pickleball (wiffle ball) + TT (tiny fast ball on a table) are out-of-distribution ? no clean trajectory yet.

319 This EMPIRICALLY validates the roadmap's binding multisport blocker (clean per-sport MIT/Apache ball detectors not identified).

320 - **PROGRESS (owner asked to complete 1/2/3 before badminton):**

321 - ? **{(1) convert-cvat DONE ? INGESTION PROVEN.** `python -m src.cli.curate convert-cvat --input <annotations.xml> --out

322 <csv> --fps 59.94` ? **pickleball 6 rallies + tabletennis 6 rallies**, sane intervals, attrs mapped. CSVs at

323 `%TEMP%\roora\{pickleball,tabletennis}\_rallies.csv` (+ extracted XMLs

324 `%TEMP%\roora\{pb,tt}\annotations.xml`). ? Rooru data

324 nits (non-blocking): pickleball `last\_shot\_outcome='0'` (off-vocab in/n_a/net/out); TT `ending\_reason` mostly 'other' with the

325 real outcome in `notes` ("out of table"/"net") ? ending_reason taxonomy mapping could improve. Both clips SHORT (97s/58s).

326 - ? **{(2) WASB-transfer test DONE (`%TEMP%\{pickleball,tabletennis}\_traj.csv`; log `~/clips/multisport\_run.log`):** WASB

327 (badminton shuttle detector) detects SOMETHING on both ? **pickleball 28% / tabletennis 39% visibility** (NOT near-zero;

328 comparable to badminton 40?67%). So it's not blind.

329 - ? **{(3) DONE ? SURPRISE RESULT (directional, n=6 each, TINY so noisy):**

```

`eval_partial_labels` on (WASB-transfer traj + Roora
330      6-rally CSV, fps 59.94, fw 1920): **TABLE TENNIS best F1 0.909** (tuned+gate2:
P1.0/R0.833, 5/6 matched, boundary err 0.49s ?
331      **WASB transfers to TT remarkably well!**); **PICKLEBALL best F1 0.308** (poor ?
wiffle ball too out-of-distribution).
332      ? **NUANCES the roadmap blocker: per-sport detector transfer is SPORT-DEPENDENT ? TT
may be servable with WASB short-term;
333      pickleball needs a proper detector.** ? BIG caveat: 6 rallies each = could be partly
luck ? owner sending more videos to
334      confirm. A "merged" prototype LOVO across badminton+pickleball+TT is the natural
next experiment once n grows per sport.
335      Roora LABELS ingest cleanly + are corpus-gold regardless. **Great honest surprise to
greet owner with.**
336
337      **Checkpoint:** 2026-06-09 (PR #61 MERGED ? roadmap SET, session wrap) . ? **Owner
agreed + merged #61 (`master 4d04a73`).
338      ALL CLEAN: master, 0 uncommitted, 0 open PRs, branches pruned.** The owned-model
**ROADMAP is now set + agreed** ?
339      authoritative doc = **`docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`** (full verified
decisions in the "(PR review)" checkpoint
340      just below: ms-swift + Qwen3-VL, TAL-head-first, conversational-QA seam-now-model-later,
multisport one-head/base, Python
341      everywhere, MIT/Apache/BSD ratified, "envision now build later"). [[session-handoff]]
rewritten to this phase.
342      - ?? **NEXT SESSION (awaiting owner greenlight ? NO build until then):** greenlight a
Phase-0 milestone (rec **M0.1
343      corpus+event-index** ? **M0.4 Gen-0 TAL harness**, both $0/local); confirm **collector
= system-of-record**; define the
344      **ship-gate F1@tIoU target** (floor = beat heuristic LOVO 0.480); **rename repo off
"badminton"**, M0.3 purge Gemini
345      training labels + WASB C3 (own-weights/clean swap). Corpus = n=6
(`output/human_lovo_manifest.json`); learned prototype
346      is the Gen-0 seed. Committed next PLATFORM slice stays the async job model (single-
tenant). Owner flagged context ~85% ? this is the wrap checkpoint.
347
348      **Checkpoint:** 2026-06-09 (PR review) . ? **OWNER REVIEWED PR #61 ? REFRAMED ROADMAP
(important checkpoint).** Owner left
349      ~12 inline comments (mostly state-updates + recommendations wanted). **Key roadmap
shifts to internalize:**
350      - **Immediate goal REFRAMED:** not "on-device model" but a **TRAINING FRAMEWORK +
HARNESS to produce a VERY HIGH
351      precision/recall model that extracts HIGHLIGHTS + RALLIES from a clip.** Quality
first. On-device = LATER (compression
352      for power users to save upload cost), definitely on roadmap but not now.
353      - **MULTISPORT FROM THE START** ? all net-separated racquet sports ? eventually all
sports; a SINGLE sports-agnostic model.
354      Collector + indexer already sport-generic by design. ? **ACTION: rename the repo off
"badminton" ASAP** (owner emphatic).
355      - **CONVERSATIONAL VIDEO-QA = north-star feature:** upload video ? ask contextual
questions ("longest rally", "make a clip
356      of all service faults", "rally >20 shots"). Needs per-video STATE/bookkeeping ? design
the framework for it. Owner asks:
357      embark now or punt + revisit after quality? (My lean: design the bookkeeping/event-
store layer now, punt the heavy
358      conversational model ? pending the strategy workflow.)
359      - **LICENSING GUARDRAIL (owner: "never corrupted")** EVERY dependency ? code, data, AND
weights ? must be usable in
360      PRIVATE PROPRIETARY COMMERCIAL software. **MIT / Apache-2.0 / BSD only.** Ratify (was
"Apache-2.0 only"). Enforce religiously.
361      - Owner wants a **coherent recommendation: local Gemma 4 vs ms-swift starting point**
(multisport, single, conversational,
362      commercial). NB: ms-swift = training FRAMEWORK, Gemma 4 = BASE model ? orthogonal;
clarify in the reply.
363      - Owner: **upgrade WASB C3 (own-weights/provenance) to a first-class action item now.**
M0.2 (grow n) = practically DONE

```

```

364         (n=6); owner will add 1 more varied multi-court-club video soon ? Phase-0 ready to
start.
365     - **LANGUAGE QUESTION (owner callout, no preference):** is Python right for the upcoming
work, or brainstorm alternatives?
366     - ? **STRATEGY WORKFLOW done (`wf_f9aa1d1c-c83`, 10 agents, verified+critiqued) ?
VERIFIED DECISIONS (now in the rewritten
367     `docs/OWNED_MODEL_IMPLEMENTATION_PLAN.md`, the authoritative roadmap; pushed to PR #61
+ a consolidated PR comment reply):**
368     - **Framework = ms-swift (Apache-2.0)** (only forge w/ native video+audio+timestamp-
grounding+video-GRPO); unsloth = 8GB-local
369     SFT/compression tool (avoid its AGPL Studio). **Base = Qwen3-VL (Apache-2.0,
timestamp emission)** primary; **Gemma-4 STAYS
370     AMBER/bench-only** (critic corrected the workflow's "clear AMBER" ? posture
unconfirmed + 60s/no-timestamp dispute);
371     InternVL3 (MIT) fallback. ms-swift?Gemma4 (framework vs base ? clarified for owner).
372     - **START WITH THE TAL HEAD, NOT THE VLM** (VLMs miss strict temporal boundaries;
boundary accuracy = the product; head
373     de-risked AUC 0.83?0.92, $0 local, doubles as teacher+baseline). Head now ? VLM
later (two-stage: VLM proposes, head refines).
374     - **Conversational-QA: build the SEAM now (per-video event store + structured text-to-
SQL + ffmpeg clip-cut), PUNT the NL
375     model** until quality clears the bar. Queries ("longest rally"/"service faults"/">20
shots") = structured-numeric ? SQL, not
376     a VLM. Reserve fault/event taxonomy in M0.1. ? real DB delta (new cols + a
`highlights` table; play_segments lacks them today).
377     - **Multisport: one shared head + one base; scale DATA + per-sport clean trajectory
detectors (the binding blocker ? none
378     identified for tennis/TT/pickleball/padel yet); WASB C3 multiplies. Single-model
temporal generalization = OPEN HYPOTHESIS
379     (RacketVision evidence REFUTED = ball-tracking not temporal). Repo rename off
"badminton" = action item.**
380     - **Language: Python everywhere** (training Python-locked; platform I/O-bound, no hot
path ? reject Go/Rust split); on-device =
381     deferred EXPORTED runtime shell (ONNX/CoreML/GGUF), doesn't dictate dev language.
Constraint NOW: **train export-clean ops**.
382     - **Honest gaps:** cost-to-Gen-2 UNQUANTIFIED (biggest risk); ship-gate "very high
P/R" target undefined (floor = beat heuristic
383     LOVO 0.480 ? set explicit F1@tIoU); per-sport detectors TBD; base PRETRAINING
provenance unauditable ? owner risk-acceptance;
384     **Gemini-silver TRAINING labels = LIVE violation ? hard-blocker purge M0.3.**
385     - **Owner principle recorded: "envision now, build later"** ? design every deferred
seam (event store, export boundary, sport-
386     agnostic head) up front; punt the feature, never the seam.
387     - ?? **NEXT:** owner reviews the reframed plan + decisions (greenlight a Phase-0
milestone ? rec M0.1 corpus+event-index then M0.4
388     harness; confirm collector=system-of-record). Then start the greenlit milestone. Repo
rename. Define the ship-gate F1@tIoU target.
389
390     **Checkpoint:** 2026-06-09 . ? **n=6: MahadevpuraSingles added while owner reviews PR
#61.** Owner moved ALL golden videos
391     + CSVs to `Annotation Setup/Collect/Golden Labelled/` (manifest label paths updated;
trajectories unaffected in `%TEMP%`).
392     New: **MahadevpuraSingles** (singles, 1080p, 29.97fps, 31 rallies; traj
`%TEMP%\mahadevpura_traj.csv`, script
393     `run_wasb_mahadevpura.sh`). **Borderline multi-court (contrast 2.2x)** ? heuristic best
0.667, learned held-out **AUC
394     0.902/F1 0.712** (clean generalization to an unseen match). **n=6 LOVO: heuristic 0.480
± 0.242 vs learned 0.583** (learned
395     holds ~+0.10 since n=3). Contrast metric now predicts heuristic F1 MONOTONICALLY across
5 matches (Mahadevpura's 2.2x?0.67
396     lands on the curve). **Decoupling headline: learned held-out AUC 0.83?0.92 regardless of
contrast; heuristic swings 0.16?0.75
397     with it.** Committed to PR #61 (n=6 doc update). ? GOTCHA: don't combine inner `&` with
run_in_background (orphans the job);
398     calibrate_local buffers output to end + is slow (~17?20min). "Ingested" = added to EVAL

```

corpus (manifest); collector

399 import-local (system-of-record) still = M0.1, plan-only pending greenlight. ?? owner reviewing PR #61 (4 plan decisions pending).

400

401 ****Checkpoint:**** 2026-06-08 (later+4, SESSION WRAP) . ? ****n=5** corpus + MULTI-COURT root cause + heuristic discovery + SINGLE PR #61 for offline review.

402 Owner tagged gBaaddy (47, 1080p, multi-court) + TestLargeVideo (6, short) ? corpus ****n=5**** (`output/human\_lovo\_manifest.json`).

403 - ****HEURISTIC DISCOVERY** (user asked "discover more heuristics like audio"): ****tested 2 label-free trajectory pre-filters**

404 for multi-court (measured on all 5): foreground spatial-density gate = ****net ?0.04**** (helps multi-court +0.03, hurts clean

405 ?0.12); track-continuity/jump filter = ****no-op****. ? simple single-cue trajectory heuristics CAN'T isolate the foreground

406 court. ****AUDIO-foreground idea PROTOTYPED** (no more videos needed) ? **WEAK ?:** ****on the 2 multi-court videos + clean**

407 control, audio onset/RMS energy AUC only ****0.53?0.65**** for in-rally vs dead-time; HF bands (8?22kHz @48kHz) raise

408 CONTRAST (Kushagra 1.1?1.7x, gBaaddy 1.3?1.9x, confirming distance-attenuation physics) but AUC stays ~0.6. Fails:

409 dead-time isn't quiet ? foreground court loud during dead-time too (footwork/talk/pickup) + gym reverb ? audio loudness

410 ? court occupancy, not rally (same wall as E1 ~2.2x). ****Audio alone ? multi-court gate.**** Remaining heuristic lever =

411 real court-region detection (player-cluster/court-line, ROADMAP #2); robust path = the learned visual model (AUC 0.85).

412 All tested + corrected in PR #61 doc + a PR comment.

413 - ****COURT-REGION DETECTION SMOKE-TESTED ? LOW HEADROOM ?:** ****only 12?14% of dead-time detections on the multi-court videos**

414 fall OUTSIDE the foreground rally region (86?88% INSIDE) ? a perfect court gate lifts contrast only ~1.6?2.2x (still

415 failing). The confusion is the foreground court's OWN between-point activity, NOT separable background courts. So court

416 detection (heuristic OR own model) won't help. ****3 heuristics now tested + FAILED** (spatial gate, audio, court

417 separability) ? all because it's a TEMPORAL rally-vs-non-rally problem in the same space; only the learned model solves

418 it (AUC 0.83). NO cheap heuristic shortcut ? the learned segmenter is the path. **** (In PR #61 doc + PR comment.)**

419 - ? ****SINGLE PR #61**** (<https://github.com/avidullu/badminton-highlight-indexer/pull/61>) carries EVERYTHING for offline

420 review: rally_seq_proto + tests, owned-model STUDY + PLAN (4 decisions pending), multivideo-lovo learnings (n=5 + multi-court

421 + heuristic discovery), CODE_MAP. Title/desc updated. Plan doc NOW committed (was held). Temp experiment scripts deleted.

422 (#60 + collector #13 merged earlier; #61 = only open PR.)

423 - ****? DEFINING FINDING** (owner's insight, data-confirmed): heuristic fails on MULTI-COURT footage. ****gBaaddy (1080p like**

424 Adarsh) heuristic best F1 ****0.277****; Kushagra 0.163 ? both have BACKGROUND GAMES. WASB tracks every shuttle incl. other

425 courts ? high dead-time visibility (gBaaddy 41%, Kushagra 57% vs clean Adarsh 18% / doubles 22%). ****"Contrast" metric****

426 (in-rally÷dead-time vis) predicts heuristic success: ?3.6x works, ?1.9x fails. Academy = multi-court = TARGET ENV ? these

427 are the most representative videos. ****Lever = foreground-court isolation**** (spatial gate to prominent court; needs real

428 court-region detection, naive centroid muddled by foreground-shuttle-resting-low). NEAR-TERM, no new model.

429 - ****Learned prototype LOVO (n=5):**** per-frame AUC 0.83?0.92 on every match; ****recovers BOTH multi-court videos****

430 (Kushagra 0.561, gBaaddy 0.574 vs heuristic 0.16/0.28). Kushagra climbs 0.049?0.417?0.561 with n. Mean 0.568 ? dipped

431 from 0.615 (n=3) ONLY due to the 6-rally TestLarge noise fold; ****excl. it = 0.639****. Per-video gap on hard videos is the signal.

432 - ****Heuristic LOVO n=5 = 0.463 ± 0.247**** (per-fold: Adarsh 0.751, doubles 0.758,

Kushagra 0.157, gBaaddy 0.295, testlarge
433 0.353). ****FELL from n=3's 0.544**** as multi-court videos entered ? the heuristic can't
represent the target env, so a more
434 representative corpus LOWERS its honest score. vs ****learned 0.568**** ? gap INVERTED
(heuristic led at n=2 +0.066; learned
435 leads at n=5 ****+0.105**** and widening). (b8jpxbryo was just slow, ~17min; completed
fine.)
436 - ****SESSION WRAP shipped ? PR #61**** (branch feat/owned-model-study-and-prototype): n=5 +
multi-court finding appended to
437 `docs/archives/quality-iterations/2026-06-multivideo-lovo/`; ****session-handoff.md**
REWRITTEN** (quality-iteration phase);
438 ****CODE_MAP updated**** (eval_partial_labels + rally_seq_proto + multi-court gotcha).
Owned model still PLAN-ONLY
439 (`docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md` DRAFT, 4 decisions pending).
440 - ?? ****NEXT (tomorrow): owner tags more varied videos**** ? re-run LOVO+prototype.
****Foreground-court isolation experiment****
441 = #1 lever. Owner reviews plan + greenlights a Phase-0 milestone (rec M0.1 corpus
spine). Commit plan doc on approval.
442
443 ****Checkpoint:**** 2026-06-08 (later+3) · ? ****n=3 CROSSOVER ? learned model OVERTAKES the**
heuristic (the data-backed green
444 light). **** User tagged a 3rd distinct match **LargeTestVideo (DOUBLES, 5.3K?1920,**
29.97fps, 49 rallies)** ? corpus now
445 n=3 (2 singles + 1 doubles, 3 cameras/frame-rates). Trajectory
`~/clips/largetestvideo\_traj.csv` (script
446 `run\_wasb\_largetestvideo.sh`). Manifest `output/human\_lovo\_manifest.json` now n=3.
447 - ****Single-video LargeTest (doubles):**** heuristic best F1 ****0.750**** (tuned, R0.918),
best-fit 0.827 ? ****doubles is**
448 heuristic-FRIENDLY; Kushagra remains the LONE heuristic failure** (not a play-type
issue).
449 - ****Honest LOVO, heuristic vs learned, n=2?n=3:**** heuristic mean held-out ****0.443?0.544**
± 0.278** (Adarsh 0.725,
450 ****Kushagra 0.151 capped at any n**, doubles 0.755). Learned prototype mean**
****0.377?0.615**** (Adarsh F1 0.701/AUC 0.878,
451 ****Kushagra 0.049?0.417**/AUC 0.850, **doubles 0.729/AUC 0.914**** trained only on the 2
singles). ****? The learned model**
452 **OVERTAKES the heuristic at n=3 (0.615 vs 0.544); at n=2 it trailed (0.377 vs 0.443).****
Crossover predicted by the plan.
453 - ****The story (honest):**** learned path SCALES with data (+0.24 for one match,
calibration gap closing ? Kushagra
454 0.049?0.417) and GENERALIZES across play type (doubles held-out AUC 0.914); heuristic
is HARD-CAPPED on footage it
455 can't represent (Kushagra ~0.15 at any n). Honest "bar to beat" is now ~0.54 (n=3),
not the fit-on-self 0.73. ?
456 Reusing every labeled video for training validated; bottleneck = data scale +
calibration (n?5), not method.
457 - ****Docs updated ? pushed to indexer PR #61**** (branch feat/owned-model-study-and-
prototype): n=3 findings + final quality
458 metrics + conclusions in `docs/archives/quality-iterations/2026-06-multivideo-lovo/`;
study + index updated. (PR #60
459 MERGED earlier; #13 MERGED.) ? Prototype speed features NOT yet fps-normalized (30 vs
60fps) yet doubles still
460 generalized (AUC 0.914) ? known refinement.
461 - ****OWNED-MODEL: PLAN-ONLY per owner**** ? `docs/OWNED\_MODEL\_IMPLEMENTATION\_PLAN.md`
written (DRAFT, uncommitted, for
462 review). User chose "just plan, don't code yet." 4 decisions pending (which Phase-0
milestone to greenlight [rec M0.1
463 corpus spine], WASB C3 stance, ratify Apache-2.0/MIT, confirm collector = corpus
system-of-record). NO owned-model
464 code until greenlight. Momentum = tag more matches (n?5) + I score each.
465 - ?? ****NEXT:**** owner reviews plan + greenlights a Phase-0 milestone; tag more matches ?
re-run LOVO+prototype (expect
466 learned mean to keep rising). gBaaddy/TestLargeVideo still untagged test videos
(trajectories ready). Commit the plan
467 doc on approval.

468
469 **Checkpoint:** 2026-06-08 (later+2) · ? **Multi-video eval corpus + owned-model
research underway.** User fully tagged

4. user (2026-06-11T12:53:06.659Z)

```
1     name: CI
2
3     on:
4       push:
5         branches: [master]
6       pull_request:
7
8     permissions:
9       contents: read
10
11    concurrency:
12      group: ci-${{ github.ref }}
13      cancel-in-progress: true
14
15    jobs:
16      # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17      # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
18      # tracked .py, including scripts/ and training/, which the backend-only
19      # py_compile smoke sweep does not reach.
20      lint:
21        name: Ruff lint
22        runs-on: ubuntu-latest
23        steps:
24          - uses: actions/checkout@v4
25          - uses: actions/setup-python@v5
26            with:
27              python-version: "3.13"
28          - name: Install ruff
29            run: python -m pip install ruff
30          - name: Ruff check
31            run: ruff check . --output-format github
32
33      # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
34      # per-module quarantine list ? ratchet by removing entries as modules clean
35      # up. Typed core deps are installed so mypy sees their real signatures
36      # (without them ignore_missing_imports would Any-stub half the checks).
37      typecheck:
38        name: Mypy (lenient baseline)
39        runs-on: ubuntu-latest
40        steps:
41          - uses: actions/checkout@v4
42          - uses: actions/setup-python@v5
43            with:
44              python-version: "3.13"
45          - name: Install mypy + typed core deps
46            run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
47genai numpy
48          - name: Mypy
49            run: mypy backend training
50
51      # Fast guard against the 7fbb0 class of breakage: a SyntaxError or
52      # module-level import error shipping while the suite stays green.
53      # Deliberately does NOT install cv2/torch/numpy ? the smoke tests must hold
54      # on machines without the GPU stack (test_import_smoke stubs what's absent).
55      import-smoke:
56        name: Syntax sweep + import smoke (no GPU stack)
57        runs-on: ubuntu-latest
58        steps:
59          - uses: actions/checkout@v4
```

```

59         - uses: actions/setup-python@v5
60         with:
61             python-version: "3.13"
62         - name: Install minimal deps
63         run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai
64         - name: Run import smoke tests
65         run: python -m pytest tests/test_import_smoke.py -v
66
67     full-suite:
68         name: Full test suite (offline, mocked)
69         runs-on: ubuntu-latest
70         steps:
71             - uses: actions/checkout@v4
72             - uses: actions/setup-python@v5
73             with:
74                 python-version: "3.13"
75                 cache: pip
76             - name: System libs for opencv-python
77             run: |
78                 sudo apt-get update
79                 sudo apt-get install -y libgl1 libglib2.0-0t64 || sudo apt-get install -y
libgl1 libglib2.0-0
80             - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
81             run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
82             - name: Install requirements
83             run: python -m pip install -r requirements.txt pytest pytest-cov httpx # httpx:
fastapi TestClient
84             # obsreport (private, soft dep) is absent here: reporting tests skip via
85             # pytest.importorskip ? run_all_tests.py prints which. The coverage floor
86             # is therefore calibrated against THIS lane's number (local runs with
87             # obsreport installed measure a little higher).
88             # Floor calibrated 2026-06-11 against this lane's measured 72% (local
89             # with obsreport: 73%). Ratchet it UP as coverage grows; never down
90             # without an owner decision.
91             - name: Run full suite (with coverage floor)
92             run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
93

```

5. user (2026-06-11T12:53:33.694Z)

```

M docs/DEFERRED_HARDENING_PLAN.md
M docs/UI_ALPHA_EXECUTION_PLAN.md
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
281a81d fix(detectors): abspath the output dir before /mnt translation
70ca6fa fix(detectors): keep ~ expandable in WSL staging/inference commands
2c241a8 refactor: collapse the hybrid-segmenter twins into a shared trajectory base
From https://github.com/avidullu/badminton-highlight-indexer
* branch          master      -> FETCH_HEAD
2c241a8 refactor: collapse the hybrid-segmenter twins into a shared trajectory base
97      fix(detectors): keep ~ expandable in WSL staging/inference commands    fix/wsl-tilde-
quoting  OPEN      2026-06-11T05:56:19Z

```

6. user (2026-06-11T12:53:35.981Z)

```

v6.0.3
v6.2.0

```


7. user (2026-06-11T12:53:58.150Z)

```
M      docs/DEFERRED_HARDENING_PLAN.md
M      docs/UI_ALPHA_EXECUTION_PLAN.md
branch 'ci/node24-action-bump' set up to track 'origin/master'.
Switched to a new branch 'ci/node24-action-bump'
  M docs/DEFERRED_HARDENING_PLAN.md
  M docs/UI_ALPHA_EXECUTION_PLAN.md
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
```

8. user (2026-06-11T12:54:04.944Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\.github\workflows\ci.yml has been updated. All occurrences were successfully replaced. (file state is current in your context ? no need to Read it back)

9. assistant (2026-06-11T12:54:13.003Z)

Now commit only the workflow file and push:

10. user (2026-06-11T12:54:19.875Z)

```
.github/workflows/ci.yml | 16 ++++++-----
1 file changed, 8 insertions(+), 8 deletions(-)
[ci/node24-action-bump 3d84869] ci: bump checkout and setup-python to Node-24 majors (v6)
1 file changed, 8 insertions(+), 8 deletions(-)
remote:
remote: Create a pull request for 'ci/node24-action-bump' on GitHub by visiting:
remote:      https://github.com/avidullu/badminton-highlight-indexer/pull/new/ci/node24-action-
bump
remote:
branch 'ci/node24-action-bump' set up to track 'origin/ci/node24-action-bump'.
To https://github.com/avidullu/badminton-highlight-indexer.git
 * [new branch]      ci/node24-action-bump -> ci/node24-action-bump
```

11. user (2026-06-11T12:54:36.318Z)

Warning: 8 uncommitted changes
<https://github.com/avidullu/badminton-highlight-indexer/pull/98>

12. assistant (2026-06-11T12:54:44.360Z)

[PR #98](<https://github.com/avidullu/badminton-highlight-indexer/pull/98>) is open. Now watching the four CI lanes until the rollout resolves:

13. user (2026-06-11T12:54:46.506Z)

Command running in background with ID: b2dljbvtd. Output is being written to:
C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-indexer\386fd981-1117-4be6-8bd0-5ef6dc747c51\tasks\b2dljbvtd.output. You will be notified when it completes. To check interim output, use Read on that file path.

14. user (2026-06-11T12:54:58.844Z)

```
1      Refreshing checks status every 20 seconds. Press Ctrl+C to quit.
2
3      Full test suite (offline, mocked)          pending 0
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
4      Mypy (lenient baseline) pending 0          https://github.com/avidullu/badminton-highlight-
```

```
indexer/actions/runs/27348227476/job/80802615837
5      Ruff lint      pending 0      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615773
6      Syntax sweep + import smoke (no GPU stack)      pending 0
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
7
```

15. user (2026-06-11T12:55:05.667Z)

```
<tool_use_error>InputValidationError: Monitor failed due to the following issues:
The required parameter `description` is missing
The required parameter `command` is missing
An unexpected parameter `task_id` was provided
```

This tool's schema was not sent to the API ? it was not in the discovered-tool set derived from message history. Without the schema in your prompt, typed parameters (arrays, numbers, booleans) get emitted as strings and the client-side parser rejects them. Load the tool first: call ToolSearch with query "select:Monitor", then retry this call. For reference, this tool's input schema is: {"\$schema":"https://json-schema.org/draft/2020-12/schema","type":"object","properties":{"description":{"description":"Short human-readable description of what you are monitoring (shown in notifications).","type":"string"},"timeout_ms":{"description":"Kill the monitor after this deadline. Default 300000ms, max 360000ms. Ignored when persistent is true.","default":300000,"type":"number","minimum":100},"persistent":{"description":"Run for the lifetime of the session (no timeout). Use for session-length watches like PR monitoring or log tails. Stop with TaskStop.","default":false,"type":"boolean"},"command":{"description":"Shell command or script. Each stdout line is an event; exit ends the watch.","type":"string"},"required":["description","timeout_ms","persistent","command"],"additionalProperties":false}</tool_use_error>

16. assistant (2026-06-11T12:55:14.690Z)

The background watch will re-invoke me when the checks resolve ? no separate monitor needed. Standing by for the CI rollout on [PR #98](https://github.com/avidullu/badminton-highlight-indexer/pull/98).

17. user (2026-06-11T12:56:31.522Z)

```
<task-notification>
<task-id>b2dljbvtd</task-id>
<tool-use-id>toolu_0176LoqRmP8cVDSRUdxdmhM</tool-use-id>
<output-file>C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-
highlight-indexer\386fd981-1117-4be6-8bd0-5ef6dc747c51\tasks\b2dljbvtd.output</output-file>
<status>completed</status>
<summary>Background command "Watch PR #98 checks until completion, then print rollout" completed
(exit code 0)</summary>
</task-notification>
```

18. user (2026-06-11T12:56:37.271Z)

```
8
9      Ruff lint      pass    11s      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615773
10     Mypy (lenient baseline) pending 0      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615837
11     Syntax sweep + import smoke (no GPU stack)      pass    15s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
12     Full test suite (offline, mocked)      pending 0
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
13     Refreshing checks status every 20 seconds. Press Ctrl+C to quit.
14
15     Mypy (lenient baseline) pass    29s      https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615837
16     Ruff lint      pass    11s      https://github.com/avidullu/badminton-highlight-
```

```

indexer/actions/runs/27348227476/job/80802615773
17 Syntax sweep + import smoke (no GPU stack) pass 15s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
18 Full test suite (offline, mocked) pending 0
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
19 Refreshing checks status every 20 seconds. Press Ctrl+C to quit.
20
21 Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615837
22 Ruff lint pass 11s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615773
23 Syntax sweep + import smoke (no GPU stack) pass 15s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
24 Full test suite (offline, mocked) pending 0
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
25 Refreshing checks status every 20 seconds. Press Ctrl+C to quit.
26
27 Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615837
28 Ruff lint pass 11s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615773
29 Syntax sweep + import smoke (no GPU stack) pass 15s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
30 Full test suite (offline, mocked) pending 0
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
31 Full test suite (offline, mocked) pass 1m27s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
32 Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615837
33 Ruff lint pass 11s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615773
34 Syntax sweep + import smoke (no GPU stack) pass 15s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
35 Full test suite (offline, mocked) pass 1m27s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615740
36 Mypy (lenient baseline) pass 29s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615837
37 Ruff lint pass 11s https://github.com/avidullu/badminton-highlight-
indexer/actions/runs/27348227476/job/80802615773
38 Syntax sweep + import smoke (no GPU stack) pass 15s
https://github.com/avidullu/badminton-highlight-indexer/actions/runs/27348227476/job/80802615851
39 Ruff lint: SUCCESS
40 Mypy (lenient baseline): SUCCESS
41 Syntax sweep + import smoke (no GPU stack): SUCCESS
42 Full test suite (offline, mocked): SUCCESS
43

```

19. assistant (2026-06-11T12:56:50.956Z)

All four lanes green ? Ruff lint (11s), Mypy (29s), import-smoke (15s), full suite with coverage (1m27s), rollup SUCCESS on every check. Merging per repo convention (squash, verified-green):

20. user (2026-06-11T12:57:00.199Z)

From https://github.com/avidullu/badminton-highlight-indexer

```

* branch      master      -> FETCH_HEAD
   2c241a8..7717e15 master  -> origin/master

```

Updating 2c241a8..7717e15

Fast-forward

```

.github/workflows/ci.yml | 16 ++++++-----
1 file changed, 8 insertions(+), 8 deletions(-)

```

M docs/DEFERRED_HARDENING_PLAN.md

M docs/UI_ALPHA_EXECUTION_PLAN.md

Your branch is up to date with 'origin/master'.

Already on 'master'

From <https://github.com/avidullu/badminton-highlight-indexer>
* branch master -> FETCH_HEAD
Already up to date.
7717e15 ci: bump checkout and setup-python to Node-24 majors (v6) (#98)
2c241a8 refactor: collapse the hybrid-segmenter twins into a shared trajectory base

21. user (2026-06-11T12:57:09.491Z)

```
M      docs/DEFERRED_HARDENING_PLAN.md
M      docs/UI_ALPHA_EXECUTION_PLAN.md
Your branch is up to date with 'origin/fix/wsl-tilde-quoting'.
Switched to branch 'fix/wsl-tilde-quoting'
  M docs/DEFERRED_HARDENING_PLAN.md
  M docs/UI_ALPHA_EXECUTION_PLAN.md
?? artifacts/
?? scratch/HIGH_RISK_REMEDIATION_PLAN.md
?? scratch/PR1_BODY.md
?? scratch/PR1_HANDOFF.md
?? scratch/PR2_deferred_hardening_status.patch
?? scratch/PR88_FIX.md
   ci/ruff-lint-lane
* fix/wsl-tilde-quoting
  master
```

22. assistant (2026-06-11T12:57:21.703Z)

Original state fully restored. Now updating the current-state memory with a micro-checkpoint per the working agreement:

23. user (2026-06-11T12:57:27.599Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

24. user (2026-06-11T12:57:34.567Z)

```
1      ---
2      name: current-state
3      description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
only on request). What just landed, what's next, open decisions."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
8      ---
9
10     # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
12     **Checkpoint:** 2026-06-11 (8?9 QUALITY SWEEP IN PROGRESS ? owner approved 8 AIs via
question UI; AI-8 packaging NOT
13     selected). **MERGED on green CI: #90** (hygiene: .coveragerc repointed post-#40 +
TEST_RUN_SUMMARY ? docs/archives),
14     **#91** (pydantic .dict()?model_dump, 7 sites, warnings 15?0), **#92** (? RUFF LANE:
ruff.toml E/F/W/B py38 +
15     fix 194 violations + `lint` CI job; **DELETED scripts/generate_khelacharya_flyer.py ?
never compiled, line 39 was a
```

25. user (2026-06-11T12:57:45.444Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\current-state.md has been updated successfully. (file state is current in your context ? no need to Read it back)

26. assistant (2026-06-11T12:57:59.451Z)

Done ? [PR #98](https://github.com/avidullu/badminton-highlight-indexer/pull/98) is squash-merged on a fully green rollout; master is now `7717e15`.

A few specifics worth knowing:

- **The latest majors turned out to be v6 for both actions**, not the v5/v6 split you guessed. I verified via the GitHub API: `actions/checkout` latest release is v6.0.3 and `actions/setup-python` latest is v6.2.0 ? both Node 24 lines. All four jobs in [ci.yml](.github/workflows/ci.yml) were bumped (8 lines total; no `with:` inputs needed changes ? `python-version` and `cache: pip` work unchanged on the new majors).
- **All four lanes passed on the PR**: Ruff lint (11s), Mypy (29s), import-smoke (15s), and the full suite with coverage floor (1m27s). The merge happened only after the statusCheckRollup showed SUCCESS on every check, per repo convention.
- **Your in-flight work was untouched**: the session started on `fix/wsl-tilde-quoting` (open [PR #97](https://github.com/avidullu/badminton-highlight-indexer/pull/97)) with the two uncommitted PR-2 doc edits in the tree. I branched off `origin/master`, committed only the workflow file, and restored the tree to `fix/wsl-tilde-quoting` afterward ? the doc edits and `scratch/`/`artifacts/` carryovers are all still in place.
- The current-state memory has a new checkpoint recording the merge.

One observation: a local branch `ci/ruff-lint-lane` still exists (leftover from the #92 era, already merged). I left it alone since branch pruning was done deliberately last session ? delete it whenever convenient.