

claude-windows / c7679509-ee2b-4443-b240-0b3b1b1a7291

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291.jsonl`  
- SHA-256: `7de442af84eedc5263001cd3e3ff7c94cc5f20a4123408edd6326225d9039a36`  
- Source modified: `2026-06-21T21:55:26+00:00`  
- Imported at: `2026-07-05T16:48:19+00:00`  
- project: `C--Users-avidu-Projects-badminton-highlight-indexer`  
- session_id: `c7679509-ee2b-4443-b240-0b3b1b1a7291`
```

Transcript

1. user (2026-06-21T13:28:24.113Z)

Hello. Can you sync to head and read these repo specially the highlight indexer and understand the current lay of the repo to propose a way of baking this into a python app which can be shipped with batteries included (o.i.e. spins off a web service to run the inference and training along with the rally annotation if needed both locally and at cloud services) and a user can download it and run the inference and training on their copy of the application without needing to source the repo. Make a plan and do exhaustive research to refine it and also use the current state docs to see the e2e CUIs completed which need to be preserved in this app mode. Can you draw out a detailed plan for this?

2. assistant (2026-06-21T13:28:58.222Z)

I'll start by following the repo-freshness convention (pull before reading) and grounding myself in the current state, then launch exhaustive research to build the plan. Let me begin with recon.

3. user (2026-06-21T13:29:04.399Z)

```
=== BRANCH ===  
feat/serving-m9-auth  
=== STATUS ===  
  M backend/config/models.py  
  M pyproject.toml  
?? artifacts/  
?? backend/api/auth.py  
?? docs/KHELSUTRA_PRODUCT_OVERVIEW.md  
=== FETCH ===  
=== LOCAL vs ORIGIN/master ===  
43eccc1 feat(serving): M8 ? per-job cost ledger + versioned pricebook + /cost API (default-OFF)  
(#288)  
43eccc1 feat(serving): M8 ? per-job cost ledger + versioned pricebook + /cost API (default-OFF)  
(#288)  
=== AHEAD/BEHIND of origin/master ===  
0          0
```

4. assistant (2026-06-21T13:29:15.622Z)

Branch is even with `origin/master` (no divergence). There's uncommitted WIP (M9 auth: `backend/api/auth.py`, `models.py`, plus a new `KHELSUTRA\_PRODUCT\_OVERVIEW.md`) ? likely the owner's or a sibling session's in-progress work. I'll treat it read-only and not clobber it. Let me pull siblings and read the key state docs in parallel.

5. user (2026-06-21T13:29:23.458Z)

```
=== rally-annotator ===
```

```

branch: main
head: d5d59f2 v1.6.4: guard unsaved rallies from silent loss + flag edit mode on Mark buttons
(#13)
dirty: 0 files
=== sports-data-collector ===
branch: master
head: 031e410 feat(annotation): richer Roora schema - court + confidence + capture report +
IAA (#29)
dirty: 1 files
=== khelsutra ===
branch: master
head: 4525c5d feat(web): B-P2 ? in-UI bucket ingest (point at a gs:// video ? process ? load)
(#56)
dirty: 2 files
=== sports-obsreport ===
branch: main
head: 1f72c1c feat: configurable highlights item_noun (v0.1.3)
dirty: 0 files

```

6. user (2026-06-21T13:29:23.526Z)

```

1      ---
2      name: session-handoff
3      description: "Living \"you are here\" handoff ? current state, next steps, open
decisions, ramp-up kit. AUTO-REFRESHED at meaningful checkpoints (PR merged, key decision, phase
shift, session wrap). \"update the handoff\" forces a full rewrite."
4      metadata:
5          node_type: memory
6          type: project
7          originSessionId: 45ea3a91-654d-4972-9cd3-65f35db54caf
8      ---
9
10     # Session handoff ? badminton-highlight-indexer (rename PENDING) + siblings: sports-
data-collector, rally-annotator, sports-obsreport
11
12     **Last updated:** 2026-06-21 (engine serving **M1?M6 + M8 shipped** ? M5/M6/M8 this
session as default-OFF, adversarial-reviewed PRs; M9-auth/M11/M12 + Ops-Agent are the remaining
authorized batch; M7 real-run + counsel/prices/OAuth = owner residual). **Repo is under the
`Khelsutra` GitHub org** (`Khelsutra/badminton-highlight-indexer`).
13
14     > **Read order:** this file ? [[current-state]] top (dense per-checkpoint detail) ?
`docs/NEXT_STEPS.md` (the prioritized cross-track plan, refreshed 2026-06-21) ?
`docs/DOC_STATUS.md` (doc lifecycle/hygiene).
15
16     ## ? You are here (2026-06-21)
17     `origin/master` = `43eccc1` . **0 open PRs** . **no GCP VMs (audited $0)** . ~1186 tests
green.
18     **? ENGINE SERVING BUILD: M1?M4 (config/input/output/device SEAM layer) ? + M5 ? + M6 ?
+ M8 ? SHIPPED.** **?? THIS SESSION (2026-06-21, 4 more PRs, owner BATCH-AUTHORIZED the default-
OFF code path):** **#285** (docs: $8 picks D16/D17 + per-video P3 reconcile) . **#286 M5**
(`backend/config/startup.py` per-profile fail/warn-fast checks + the **L4 profile-parity**
proof; default-OFF, worker stays torch-free) . **#287 M6** (`backend/compute/` ComputeBackend
registry ? `CloudRunCompute` dispatch a Cloud Run **Job** + bucket-poll a done-marker;
`deploy/cloudrun/{Dockerfile,README}`; mocked, default-OFF) . **#288 M8** (per-job **cost
ledger** + v6 migration + versioned `pricebook` seeded $0 + `/cost` API + `BillingConfig`;
default-OFF). Each isolated, **adversarial-reviewed (a workflow per PR)** , full-suite + cross-OS
CI green, tracker+changelog in-PR. Owner picks: blanket-YES to write+merge default-OFF code,
merge-each-on-green, **GCP spend OK to $100/?10k** for real M6/M7 verification, design defaults
(Cloud Run Jobs . libx264 encode+NVDEC decode . hard-reject <2GB). **M1?M4 (prior session)
detail below.**
19     - **M1 ([#279](https://github.com/Khelsutra/badminton-highlight-indexer/pull/279),
`ce22791`):** `deployment.profile`+`compute_target` on `AppConfig` + `backend/config/presets.py`
(PROFILE_PRESETS + `suggest_profile`) + loader preset/precedence
(defaults<preset<config.json<env via `_deep_merge`; no-profile path = exact old shallow merge) +
auto-detect-as-SUGGESTION (D15). Pure/additive, profile='' == today. 3 minor review findings

```

folded.

20 - **M2 ([#280](https://github.com/Khelsutra/badminton-highlight-indexer/pull/280),
`e111add`):** `StorageConfig.input\_backend`/`input\_root` +
`backend/storage/{ref.resolve_input_uri,
__init__.resolve_input_ref/acquire_local_input/input_is_local}` + `main.py` API/CLI wired
(acquire local; keep source URI for DB/report). **local = identity passthrough (byte-for-byte
unchanged)**; bucket/Drive fetched to scratch (mirrors worker `cmd\_infer`). Unknown scheme?clean
400 (was a 500 regression ? caught by review); `local://` accepted on resolved key; hosted-mode
rejects non-local; OS-independent anchored-path detection (os.path.isabs differs Win+py3.13). 1
major + 6 minor review findings folded. Suite 1105 green, cov 84.84%.

21 - **M3 ([#282](https://github.com/Khelsutra/badminton-highlight-indexer/pull/282),
`2afd54f`):** configurable output/scratch base ? `backend/utils/paths.output\_base(config)`
(returns `output.output\_dir`, default `"output"`); routed all 6 hardcoded
`os.path.join("output", ?)` scratch sites in `trajectory\_hybrid`/`yolo\_hybrid`/`gemini` through
it (create+cleanup same var). **DEFAULT `"output"` ? byte-identical** (golden + cross-OS green);
a custom `--output-dir` now actually moves scratch (latent bug fixed). Default-OFF ? no Launch
Report. `tests/test\_output\_base.py` + a yolo call-site capture; clean-bill focused review (yolo-
coverage minor folded). Suite 1114 green.

22 - **M4 ([#283](https://github.com/Khelsutra/badminton-highlight-indexer/pull/283),
`cee4baa`):** `DeviceContext` (`backend/pipeline/detectors/device.py`) + native WASB **CPU-
fallback** via a new `device="auto"` (cuda-if-available-else-cpu). **Explicit `"cuda"` stays
STRICT** (healthcheck still hard-fails without a GPU ? no silent slow-CPU downgrade; "cuda
default unchanged"). CUDA probe **cached** (healthcheck?run_predict reuse);
`\_effective\_device()` resolves before argv so "auto" NEVER reaches `wasb\_infer.py`; fail-fast
on a typo'd device. `wasb\_infer.py` already supports cpu ? no change there. Device-selection
tested offline (mocked subprocess); **real CPU-inference run = owner/GPU-verified.**
`tests/test\_device\_context.py` + native-runner additions; clean-bill focused review (2 nits
folded). Suite 1126 green, cov 84.96%.

23 **? NEXT = the REMAINING authorized batch (default-OFF code, ready to build; blueprints
in [[current-state]] top):** **M9-auth** (add `PyJWT[crypto]`; Cf-Access JWT verify dep ?
`user\_id` + per-tenant scoping on the v4 cols + tighten CORS; behind
`security.edge\_auth\_enabled=False`; offline RSA/fake-JWKS tests) . **M9-consent** inert NULL
cols in a **v7** migration (? M8 took v6 ? next bump is v7; ToS/DPA text = owner/counsel) .
M11 flywheel plumbing (`flywheel.capture/shadow\_eval/promote\_stub` reusing
`workspace.py`+`experiment.compare`+`promotion.py`, `flywheel.\*`=OFF) . **M12**
`GDriveApiStorage`+`gdrive-api://` (fake-Drive DI tests; keep mount-only `gdrive://`) . **Ops-
Agent codify** into `deploy/gcp/provision.sh`. **? owner-only residual:** the real M5 truly-
local runlog (template pre-staged) . GCP spend for the real M6/M7 run within the \$100 cap (build
image ? L4 ? DEPLOY_REPORT ? DELETE) . counsel ToS/DPA (M9-consent) . real GCP prices+FX + M8
usage-emission follow-up . Google OAuth app (M12) . CF team-domain/AUD+ingress-lock . repo
rename . DO droplet/token retire. **Session wrapped at a clean 4-PR checkpoint.**

24

25 Prior session delivered the **engine compute/serving + doc-hygiene** track (10 PRs
merged: #263?#267, #270, #272?#274, #276):

26 - **M2 FINALIZED on a real GPU** ? `worker train --trajectory-source detector` on a GCP
L4 (Singapore; Mumbai was stocked-out) ? Gen-0 bundle `gs://khelsutra-rally-
corpus/weights/0.1.0-l4/` (honest **n=2 plumbing** result, not a quality model). Both VMs
deleted.

27 - **DigitalOcean DROPPED for compute ? GCP is the sole stack** (ADR-013; DO GPU not
enabled/credit-funded, Paperspace subscription-gated). [[compute-stack-gcp-only]] . [[gcp-vm-
always-delete]] . [[gcp-ops-agent-enable]].

28 - **DECOUPLED_COMPUTE delivered + ARCHIVED** ?
`docs/archives/past\_projects/DECOUPLED\_COMPUTE/`.

29 - **NEW tracked subproject: `docs/COMPUTE\_DECOUPLED\_SERVING/` + ADR-014** ? one pure
core (DETECT?WINDOW?STITCH, never branches on profile) + `ComputeBackend`/`StorageBackend`
profiles (**truly-local**, **cloud-serving** [Cloud Run GPU runs detect AND stitch, metered],
cpu-mock). **? AWAITING OWNER M0 DESIGN SIGN-OFF before any code.**

30 - **Doc governance:** `docs/DOC\_STATUS.md` register + **archival due-diligence**
(honestly-update?then-archive; codified in CLAUDE.md; [[doc-status-register]]); per-video docs
reconciled; **data layer** carved into `docs/data\_pipeline/` (no-ML), field-
ops/PMF?`archives/field-pmf/`, Roora?`archives/roora-vendor/`.

31

32 ? **Shared checkout** ? sibling sessions commit + flip branches here mid-session (hit
twice this session). **Branch from `origin/master` explicitly, re-verify HEAD before commit,
stage explicit paths, never `git add -A`.** [[gotcha-shared-checkout-branch-collisions]]. ?

```

**`sed -i` over a glob churns CRLF on all matched files** ? separate real changes with `git diff
--ignore-cr-at-eol` and revert the noise.
33
34     ## ? Next steps (resume here)
35     1. **? M0?M4 SHIPPED (#279/#280/#282/#283) ? the config/input/output/device SEAM LAYER
is COMPLETE ? NOW M5+ (mostly owner-gated).** Decisions locked in README **$2 D7?D15** (D7
scale-to-zero · D8 versioned DB pricebook USD/INR · D9 Cloudflare Access · D10 data-for-reels
trade adults-only · D11 Gemini-ON-until-floor + D12/M11 flywheel · D13 user-switchable profiles
· D14 NORTHSTAR gdrive-api OAuth = M12 · D15 suggest?confirm). **Engine build sequence:** ~~M1
deployment.profile~~ ? ~~M2 input_backend~~ ? ~~M3 output-base~~ ? ~~M4 DeviceContext + CPU-
fallback~~ **ALL DONE.** **? M5** = truly-local profile e2e (preset + startup checks + first
committed runlog) ? **? owner real-GPU** (needs a GPU box / owner's WSL machine to run
detect+stitch locally, then commit the runlog). **M6/M7** Cloud Run dispatch + container image
(? owner GCP spend), **M8** cost ledger + pricebook, **M9** edge-auth (Cf-Access), **M11** data-
flywheel harness, **M12** gdrive-api OAuth. **Per-video-workspace P3 (`output/chunks_`) is now
subsumed by M3 ? reconcile/close it.** **Claude-doable WITHOUT owner gating:** the M5
preset+startup-check code (the actual run is owner's); the M8 cost-ledger DB migration +
pricebook + aggregation (pure code, owner does the real-cost calibration); the M11 harness
plumbing. M6/M7/M9/M12 need owner cloud/counsel sign-off.
36     2. **Per-video workspace P1?P6** (#264 shipped P0): P1 v6 DB migration `workspace_uri` ?
P2 worker ? P3 `output/chunks_` de-hardcode ? ? (Claude-doable, default-OFF).
37     3. **Ops-Agent codification into `deploy/gcp`** (auto-enable observability on every VM)
? small, [[gcp-ops-agent-enable]].
38     4. **[owner]** Grow the golden corpus (highest-leverage) · set ship-gate target #172 ·
full-corpus training (n?5) for real weights.
39     5. Full prioritized P0/P1/P2 list = `docs/NEXT_STEPS.md` top.
40
41     ## ? Open decisions / owner gates
42     - **M0 design sign-off** (the blocker above) + the COMPUTE_DECOUPLED_SERVING $8
questions.
43     - Doc-hygiene leftovers (`DOC_STATUS.md` $4): confirm the 3 COMPLETED top-level pointer-
stubs; `HOSTING_PLAN`/`UI_PROPOSAL` currency.
44     - Housekeeping: retire the shared D0 droplet `claude-do-rally-test` (~$24/mo) + the D0
token; rename repo off "badminton".
45
46     ## ? Ramp-up kit
47     - **Memory:** [[current-state]] (read first) · [[compute-stack-gcp-only]] · [[gcp-vm-
always-delete]] · [[gcp-ops-agent-enable]] · [[doc-status-register]] · [[gotcha-shared-checkout-
branch-collisions]].
48     - **Docs:**
`docs/COMPUTE_DECOUPLED_SERVING/{README,TESTING_STRATEGY,architecture.svg,runlogs/}` ·
ADR-012/013/**014** in `docs/archives/decisions/DECISIONS.md` · `docs/NEXT_STEPS.md` ·
`docs/DOC_STATUS.md` · `docs/PER_VIDEO_WORKSPACE.md` · `deploy/gcp/README.md` (GPU runbook;
teardown defaults to `delete`).
49     - **Code anchors (the pure core the serving project builds on):**
`backend/pipeline/detectors/base.py` (DetectorRunner CSV contract) ·
`tracknet_runner.py::trajectory_to_action_windows` ·
`backend/pipeline/stitcher/compiler.py::compile_highlights` ·
`backend/pipeline/segmenters/trajectory_hybrid.py::_resolve_runner` (the compute seam) ·
`backend/storage/*` · `backend/worker/___main___py` · `backend/utils/workspace.py`.
50
51     ## ? Parallel track (NOT this session ? sibling sessions)
52     **KhelSutra brand-site + product UX/Hindi-Kannada localization** (repo
`Khelsutra/khelsutra` + #275 here) ? [[next-session-ux-localization]] / [[khelsutra-domain-
launch]]. The **rally-quality / owned-model** track is live too (R-series,
`RALLY_QUALITY_RESEARCH.md`; gated on corpus growth). These are background to the
compute/serving resume above.
53

```

7. user (2026-06-21T13:29:23.527Z)

```

1     ---
2     name: next-session-ux-localization
3     description: Consolidated handoff/plan for the NEXT session ? focus SOLELY on (A) making
the UI super user-friendly for non-tech-savvy users and (B) localizing in Hindi + Kannada.

```

```

Grounded in existing repo docs (owner directive 2026-06-21).
4     metadata:
5         node_type: memory
6         type: project
7         originSessionId: 577fd4df-ed3f-43d3-854d-d63d732e2e07
8     ---
9
10    **Owner directive (2026-06-21, session wrap):** the next session focuses **SOLELY on
making the product super user-friendly (non-tech-savvy) AND localized (Hindi + Kannada)**.
*Research first using the existing repo docs, THEN craft the plan* (don't reinvent ? there's
prior design). The end goal: **a non-tech-savvy coach/parent reads the page, confidently sets up
recording, runs the tool, and gets highlights ? in their own language.**
11
12    ## ? NOW ISOLATED INTO A TRACKED PROJECT (2026-06-21 ? done)
13    The work is isolated as ONE umbrella tracked project:
**`khelsutra/docs/FRIENDLY_AND_MULTILINGUAL.md`**
14    (workstream A = non-tech UX, B = kn/hi l10n; $7 P-tracker = source of truth;
**DoD/archive trigger = en+kn+hi
15    render on BOTH `web/` SPA + `site/` marketing**). Inbound refs point at it: khelsutra
`docs/README.md` +
16    `NEXT_STEPS.md`; engine `I18N_PLAN.md` + `UI_PROPOSAL.md` banners + `DOC_STATUS.md`
$3a/$3b rows. **Branches
17    OPEN: khelsutra docs [#41], engine pointers [#275], B1 scaffold [#42].** Owner locked D6
(en+kn?hi) · D7
18    (OFL Noto self-host) · D8 (SPA-first) · D9 (catalogs in khelsutra); only D-BACKEND-DEPTH
deferred.
19    **B1 DONE + CI GREEN + verified live** (react-i18next + ICU + detector + `en` catalog +
language switcher +
20    the SPA's first jsdom render tests; 79 tests green). **? NEXT (work the doc's $7, don't
re-plan): B2** (full
21    ~50-string extraction + key-parity/no-bare-strings CI) ? **B3** (site `t()` shim) ?
**B4/B5** (Noto fonts +
22    kn catalog via Gemini-draft?human-review).
23
24    ## Clean slate at handoff (2026-06-21)
25    Both repos on `master`, **0 open PRs**, no orphaned worktrees, **no GCP VMs** (no
billing). khelsutra master `5de981f`; engine (`Khelsutra/badminton-highlight-indexer`) master
`4134721` (DECOUPLED_COMPUTE archived as delivered). Engine local checkout is the **concurrent
session's** branch ? don't disturb it; use an isolated `git worktree` off `origin/master` for
any engine work (the #116/#117 collision rule).
26
27    ## ? Ramp-up kit (READ THESE before crafting the plan ? the owner's "use the docs" ask)
28    **Localization (the pillar):** **`engine docs/I18N_PLAN.md`** ? Kannada-first?Hindi-
next, **by design** (no hardcoded strings; `t("key")` + ICU MessageFormat plurals; one `en`
source catalog + per-locale siblings; locale-negotiation order = persisted pref ?
switcher/?`lang=` ? `Accept-Language` ? `en` fallback; **Gemini-DRAFT?human-REVIEW** [inference,
never training ? guardrail-OK]; CI **key-parity** + **no-bare-strings** lint; **Noto Sans
Kannada/Devanagari**, OFL font decision). Phases P0?(doc)?P1(static-UI retrofit)?P2(backend
text)?P3(**SPA react-i18next**)?P4(conversational/coach-mode).
29    **UX + personas (the pillar):** **`engine docs/UI_PROPOSAL.md`** $2 personas
(Coach/academy = primary B2B India, Kannada-comfortable, records on static cameras; Enthusiast
B2C) + the "teach recording BEFORE upload / zero hand-holding" principle + the 6 alpha screens
(A1 Matches · A2 Upload+recording-checklist · A3 Job progress · A4 Rally review ? · A5
Highlights · A6 Report).
30    **Non-tech-savvy reality:** **`engine docs/PMF_FIELD_SIGNALS.md`** +
**`docs/PRODUCT_FIELD_ASSOCIATE_FLYER.md`** (field associate works in **Kannada/Hindi/English**;
coaches/players/**parents**; "simple summary"; **WhatsApp-native**; possible B2C-to-parents
pivot). **`docs/PERSONALIZATION_PLAN.md`** (coach overlay).
**`docs/KHELSUTRA_PRODUCT_OVERVIEW.md`** (who it's for).
31    **Current product surfaces (what to localize + simplify):** khelsutra **`web/` SPA**
(per-rally operator console ? `docs/PER_RALLY_SELECTION.md`) + **`site/`** marketing (homepage,
`/capture` checklist, `/own-model`, request-access form, post-signup reveal). `web/README.md`.
32
33    ## ? THE adaptation the next session must make
34    `I18N_PLAN.md` was written for the **OLD static UI** (`backend/api/static/*`). The

```

product is now the khelsutra **web/ SPA** (Vite/React/TS/Tailwind) + the static marketing **site/** (vanilla HTML). So adapt the design's framework-agnostic foundation to:

35 - **web/ SPA ? react-i18next** (the doc's P3) consuming `locales/{en,kn,hi}/*.json` catalogs (ICU plurals); a language switcher + persistence; localize the per-rally UI (query examples/pills, rally cards, footer, build/compile/download, errors). Honest scope still applies (duration-first; shot-count experimental).

36 - **Marketing site/** (vanilla) ? a tiny `t(key,vars)` shim + `data-i18n` over `index.html/capture.html/own-model.html` (no framework), sharing the SAME catalog format. The form + reveal copy + the recording gist localize too.

37 - **Backend (engine)** ? P2 (errors/report `customer_*` as keys+params) is owner-gated + a separate worktree PR.

38

39 **## Two threads to plan** (one tracked project doc each, per `[[working-agreement-tracked-project-docs]]`)

40 ***(A) Non-tech-savvy UX** ? research ? plan. Likely levers (validate against the docs + personas): plain-language/low-jargon copy, fewer steps + bigger touch targets, **mobile-first** (coaches use phones), visual/iconographic over text, **recording-education-first** (? already shipped the homepage "Record it right" gist + `/capture`), WhatsApp-friendly share, a guided "upload ? wait ? here's your reel" happy path with zero knobs. The per-rally query bar already leads with **duration** and flags shot-count "experimental" ? keep simplifying.

41 ***(B) Hindi + Kannada localization** ? adapt `I18N_PLAN.md` to `web/` + `site/`; stand up the catalog + negotiation + a first `kn` (then `hi`) draft via `scripts/i18n_draft.py`-style Gemini-draft?owner-review; CI key-parity; Noto fonts.

42

43 **## Open OWNER decisions to resolve first** (from `I18N_PLAN §4`)

44 1. **Indic font licensing** ? recommend (a) **OFL exception for bundled fonts**, self-host Noto Sans Kannada/Devanagari (treats fonts as media, not code/weights ? also unblocks a Kannada/Hindi reel watermark via the existing `WatermarkConfig.font_path`).

45 2. **v1 locale set** ? recommend `en` + **kn** for v1, **hi** next; always-fallback to `en`.

46 3. **Backend-localization depth** ? recommend customer verdict / report `customer_*` + validation-rejection reasons first; operator/technical text stays English.

47 4. **UX scope** ? which surface first: the marketing `site/` (first-contact friction) and/or the `web/ SPA` (the operator console).

48

49 **## Already shipped this session** (DON'T redo ? build on it)

50 Per-rally MVP P1?P5 (query bar/grid/footer/build?compile?download, honest duration-first, shot-count parked experimental); engine **M1** (`/api/health` + compile `download_url`); marketing site ? post-signup reveal (tool screencast + 30s reel, **watchable**: manual close + holds-while-playing), **submit "Sending?" feedback**, **"Record it right"** homepage gist (from `VIDEO_RECORDING_GUIDELINES.md`) + `/capture`; **behavioral jsdom tests** (test-only dep; runtime stays dependency-free); worker **observability fix** (logs a 0-segment cold-start, [#271]). Promo/CUJ screencasts in `gs://khelsutra-rally-corpus/promo/`. ? Cold-start blocker still open: native-WASB `worker infer` config that reproduces the engine session's 22-rally result (see `[[current-state]]`). `[[monetization-plan]]` `[[khelsutra-domain-launch]]`

51

8. user (2026-06-21T13:29:24.021Z)

<system-reminder>[Truncated: PARTIAL view ? showing lines 1-50 of 111 total (47083 tokens, cap 25000). Call Read with offset=51 limit=50 for the next page, or Grep to find a specific section. Do NOT answer from this page alone if the answer may be further in the file.]</system-reminder>

```
1 ---
2 name: current-state
3 description: "Live micro-checkpoint state for Claude ? updated every checkpoint (not
4 only on request). What just landed, what's next, open decisions."
5 metadata:
6   node_type: memory
7   type: project
8   originSessionId: b13673cc-f087-43b1-9527-35dc0b67b301
9 ---
10 # Current state ? updated continuously (see [[working-agreement-checkpoints]])
11
```

12 > ****Scope of this file = LIVE only:**** the current handoff + ***today's*** checkpoints + open PRs/decisions. Keep it read-first-short (?1 page). Older checkpoints are archived in `[[current-state-archive]]` ? when today's checkpoints age out (next clean-slate day), move them there, don't let them pile up here.

13

14 *****? NEXT SESSION FOCUS (2026-06-21, after the big localization+dry-run session):**** SPA is fully ****en+kn+hi (DONE)****; a real ****UI?engine?reel?download CUJ ran (en+hi)****; a ****critical engine bug fixed**** (#281). *****? NEXT BUILD (owner-picked):** (1) Path B ? point at a bucket video IN THE UI.**** ? **B-P1 VERIFIED ALREADY DONE ? NO ENGINE CODE (do NOT re-scope it as a feature):**** ``POST /api/process`` has accepted ``gs://`/`gcs://`` URIs since ****M2/#280**** (submit resolves the URI + skips the local-exists-404 for non-local backends ``main.py:283``; the worker fetches to scratch ``acquire_local_input main.py:159`` ? same hash?validate?segment?DB pipeline ? queryable API-DB ``video_id``); ****TESTED**** (``tests/test_api_endpoints.py::test_process_nonlocal_uri_skips_exists_and_fetches`` + hosted-mode-reject / unknown-scheme-400 / CLI variants). Confirmed by a ****3-skeptic verification workflow (0 refuted, gaps:[])****. ****Scope pin:**** holds in DEFAULT single-user mode (hosted-mode ``allowed_video_root`` correctly ****400s**** a bucket URI ? by design). Stale "local-path-only / net-new" docs ****CORRECTED ? khelsutra [#55]**** (DRY_RUN_CUJ \$1/2/4/6/8 + LOCAL_APPLIANCE + README). ****? So Path B = ONLY B-P2 = the SPA ingest/progress screen**** (enter a ``gs://`` URI ? ``POST /api/process`` ? poll ``/api/jobs`` ? load ``result.video_id``; back-half already works) ? ****100% khelsutra `web/` SPA + i18n (en/kn/hi), ZERO engine code.**** Ready B-P2 sketch: ``DRY_RUN_CUJ.md`` \$2 (new ``web/src/{api/types,api/client,hooks/useIngestJob,components/IngestPanel}.tsx`` + ``App.tsx`` ``videoId?state handoff + `ingest.`` keys). See `[[lesson-verify-engine-surface-before-scoping]]`. **** (2) Droplet deploy (opt 3):**** deploy the FIXED engine to ``168.144.126.176`` (un-provisioned, 2-core/3.8GB) ? re-run CUJ on Linux. ****STILL OPEN (localization DoD):**** B3 (``site/`` marketing still English-only) + native-speaker review of the AI-drafted kn/hi. ****Ramp-up:**** ``khelsutra/docs/DRY_RUN_CUJ.md`` + ``FRIENDLY_AND_MULTILINGUAL.md`` \$7 + the dry-run checkpoint below. ****Clean slate (at session start):**** both repos on master, 0 open PRs, no worktrees/VMs. khelsutra master ``2421557`` · engine master ``2afd54f`` (? origin/master ``cee4baa``, #283 M4 squashed).

15

16 ****Checkpoint:**** 2026-06-21 (ENGINE SERVING M5+ ? OWNER BATCH-AUTHORIZED the independent path). Owner redirected: "first surface the owner-gated decisions so I can take them + unblock max independent PRs." Ran a 6-agent grounded gate-map ****workflow**** (``wf_4849400f``) over README \$2/\$7/\$8 + code ? a decision menu splitting every M5?M12 gate into **** (a) default-OFF code I can merge now**** vs **** (b) thin owner residual****. ****Owner answers (AskUserQuestion):**** **** (1) BLANKET YES**** ? write+merge ALL default-OFF serving code (M5,M6,M8,M9-auth,M9-consent-cols,M11-plumbing,M12), one isolated PR each; ***** SPEND OK up to \$100/?10k**** for due-diligence + ****production-grade testing/verification**** (so the real M6/M7 cloud runs are ALSO unblocked within the cap ? must respect it, DELETE resources after, confirm cmd+cost before each spend per `[[gcp-vm-always-delete]]`). **** (2) MERGE each on green CI**** (OAuth real-testing deferred to a later rigorous pass). **** (3) ALL recommended design picks**** "with generous budget in mind": ****co-located stitch**** (D4) · ****libx264 ENCODE + NVDEC DECODE-only**** (protect parity) · ****hard-reject <2GB` v1**** (config knob) · confirm ``COMPUTE_DECOUPLED_SERVING/`` name · ****M6 = Cloud Run JOBS + bucket-poll completion****. ****? PLANNED PR ORDER:**** A(doc: \$8 resolutions + P3 reconcile + tracker housekeeping) ? B(M5 preset+startup-checks+L1?L4 parity test) ? C(M6 ``backend/compute/`` registry + Dockerfile + mocked-dispatch) ? D(M8 v6 migration + pricebook[\$0 placeholder] + USD/INR aggregation + /cost API) ? E(M9-auth Cf-Access JWT verify + tenant scoping; M9-consent inert NULL cols EXTEND D's v6 migration) ? F(M11 capture?store?shadow-eval?goldenize plumbing, ``flywheel.*`=OFF) ? G(M12 `GDriveApiStorage`+`gdrive-api://`, fake-Drive tests) + Ops-Agent codify into `deploy/gcp/provision.sh`. Coordinate: D/E/M9-consent share ONE v6 migration (author in D, extend in E). **Genuinely owner-only residual:** real GPU run (M5 runlog) [or I run on cloud GPU within cap], counsel ToS/DPA (M9-consent text), Google OAuth app (M12 real), real GCP prices+FX (M8 calibrate), CF team-domain/AUD+ingress-lock, repo rename, D0 droplet/token retire. ? shared checkout ? sibling merged **#284** (`06ff4b7`, bucket-URI regression guard) mid/just-before this session; RE-FETCH origin/master before each branch. **? PROGRESS (this session):** ? **PR-A [#285] MERGED** (`fefa412`, doc: $8 D16/D17 resolutions + P3 reconcile + tracker housekeeping). ? **M5 [#286] MERGED** (`b5af66e`, all 7 CI green) ? `backend/config/startup.py` per-profile fail/warn-fast checks (cloud-serving+non-cloud-storage = the one fatal coherence error; GPU/creds = warnings) wired into server (`_enforce_startup_or_exit`?exit1) + worker (`cmd_infer`/`cmd_train`?rc2), **default-OFF + probe gated behind active profile (worker stays torch-free at profile='')**; **L4 profile-parity test** (`tests/test_profile_parity.py`: same stub trajectory?byte-identical windows+ranges across truly-local/cpu-mock/no-profile + worker-vs-in-process); 4-lens adversarial review folded (`wf_34671ac1`: MAJOR torch-free fix + minors). Runlog template pre-staged; the real-GPU truly-local **runlog = owner L5**. M5 ? in $7 tracker.`

? **M6 [#287] MERGED** (`8cdf65c`, all 7 CI green) ? `backend/compute/` ComputeBackend registry: `get\_compute\_backend(config)` returns a backend ONLY for `compute\_target=cloud\_run` (else `None` ? worker's in-process path, **default-OFF**); `CloudRunCompute` dispatches a Cloud Run **Job** (injectable `CloudRunJobClient`; real `GoogleCloudRunJobClient` lazy google-cloud-run) ? **bucket-polls** a done-marker via `execution\_policy.poll\_until`+storage (D16); container forced **inline** (`compute\_target=local\_gpu`+native, never re-dispatch); worker `--done-uri` writes a success marker. `deploy/cloudrun/{Dockerfile,README}` = L4 image (CUDA+ffmpeg+DejaVu fonts+`wasb` conda torch-1.11 env mirroring gpu_setup.sh; weights staged per WASB-C3) + M7 runbook. 4-lens review (`wf\_87b22fa4`) folded a **BLOCKER** (cloud `fetch` needs a dst ? reads silently returned `{}` on gcs/s3; tests passed only b/c local) + poll-wait test gap + `PolicyTimeout`?clean rc 4. Suite 1167 green. ? **M8 [#288] MERGED** (`43eccc1`, 7 CI green) ? `backend/billing.py` (pure `usage`+pricebook, USD source-of-truth / INR display) + **v6 migration** (NULLABLE cost cols on `jobs` + versioned `pricebook` seeded `placeholder-v0` at \$0; additive=NULL, idempotent) + `record\_job\_cost`/`get\_job\_cost`/`get\_video\_cost` + `GET /api/jobs|videos/{id}/cost` + `BillingConfig` (**default-OFF**) + gated `\_maybe\_record\_job\_cost`. 3-lens review (`wf\_498b24af`) folded a **MAJOR silent-GPU-under-bill** (an unpriced `gpu\_type` is now SURFACED ? WARNING + `unpriced\_gpu\_seconds` breakdown line, not invisible \$0) + the **honest no-op** caveat (pipeline doesn't emit `result['usage']` yet ? cost 0 until a follow-up wires the metering; documented + tested) + margin/rounding reconciliation. Suite 1186 green. **?? 4 SERVING PRs THIS SESSION, ALL MERGED: #285(docs \$8 D16/D17 + P3 reconcile) + #286(M5 startup-checks+L4 parity) + #287(M6 Cloud Run dispatch shim+image) + #288(M8 cost ledger).** Each default-OFF, isolated, adversarial-reviewed, full-suite+CI green, tracker+changelog in-PR. **? v6 is now SHIPPED (M8 took it) ? the NEXT migration is v7** (M9-consent cols + per-video P1 `workspace\_uri` ? the earlier "extend v6" plan shifts to a fresh v7 `\_migrate\_to\_v7`). **? REMAINING authorized batch (ready to build next, blueprints loaded):** **M9-auth** = add `PyJWT[crypto]`; a FastAPI dep extracts `Cf-Access-Jwt-Assertion` ? fetch+cache the team JWKS ? verify sig/exp/iss/aud ? `user\_id`; per-tenant scoping on the v4 `user\_id` cols; tighten the placeholder CORS (`main.py:40`); all behind `security.edge\_auth\_enabled=False`; offline RSA/fake-JWKS tests. **M9-consent** = inert NULL `consent\_status/is\_adult\_attested/retention\_class` in a **v7** migration (safe no-train default) ? the ToS/DPA TEXT is owner/counsel. **M11** = `flywheel.capture/shadow\_eval/promote\_stub` reusing `workspace.py::for\_video` + `backend/eval/experiment.compare`(:485) + `promotion.py::promote\_if\_served\_safe`, behind `flywheel.\*`=OFF + a faked `consent\_attested`. **M12** = `GDriveApiStorage(StorageBackend)` + a `gdrive-api://` scheme against an INJECTED fake Drive (mirror `gcs.py` DI), keep the mount-only `gdrive://` (truly-local) untouched. **Ops-Agent codify** into `deploy/gcp/provision.sh` (logging+monitoring APIs + SA metric/log-writer roles + startup-script ? [[gcp-ops-agent-enable]]). **? owner-only residual:** real GPU run ? M5 truly-local runlog (template pre-staged in `runlogs/`); GCP spend for the real M6/M7 run within the **\$100/?10k cap** (build the image per `deploy/cloudrun/README`, run on L4, capture DEPLOY_REPORT, **DELETE after**); counsel ToS/DPA (M9-consent); real GCP prices+FX + the M8 usage-emission follow-up; Google OAuth app (M12 real); CF team-domain/AUD + ingress-lock; repo rename; DO droplet/token retire. **Session wrapped at a clean 4-PR checkpoint** ? M9/M11/M12 + Ops-Agent handed off ready (per the quality bar: isolated PRs + rising coverage + adversarial review each). [[working-agreement-merging]] [[feedback-launch-quality-bar]] [[compute-stack-gcp-only]]

17

18 **Checkpoint:** 2026-06-21 (STALENESS-FIX / Path-B verify) ? owner: pick up Path B + "fix all stale docs and stop next sessions spiralling into the staleness." **? Verified-FIRST finding: B-P1 was ALREADY DONE** ? `POST /api/process` accepts `gs:///`/`gcs:///` since **M2/#280** (resolve?skip-exists-404?worker fetch?same pipeline?queryable API-DB `video\_id`) + TESTED; confirmed by inline file:line trace AND a **3-skeptic verification workflow (0/3 refuted, gaps:[])**. So the doc-scoped "build the gcs ingest endpoint" was a phantom; Path B = **SPA-only (B-P2)**. **DONE this session:** (1) **khelsutra [#55] MERGED** (master `d0461fd`) ? corrected `DRY\_RUN\_CUJ` \$1/2/4/6/\$8 + `LOCAL\_APPLIANCE` (`/api/health` EXISTS now M1/#268; bucket-URI ingest NOT a net-new endpoint; T12 engine-half ? done) + `README`; added the anti-spiral convention "re-verify engine-surface `[gap]`/`net-new` claims at a cited file:line ? the engine ships PRs between sessions." (2) **Memory steering corrected** (this NEXT-SESSION-FOCUS block + the dry-run checkpoint were stale "local-path-only") + new lesson [[lesson-verify-engine-surface-before-scoping]] + MEMORY.md index. (3) **engine [#284] MERGED** (master `06ff4b7`, all CI green incl. full suite 2m32s) ? regression GUARD `test\_process\_bucket\_uri\_yields\_queryable\_video\_id` (parametrized gs://+gcs://) pinning `video\_id`==MD5(fetched bytes) (not the URI) + `POST /api/query` queryability (the SPA contract; previously only LOCAL had it); built+removed in an isolated worktree off `origin/master` (engine checkout left on its concurrent-session branch). **? NEXT: B-P2** = the SPA ingest/progress screen (enter gs:// URI ? /api/process ? poll /api/jobs ? load video_id; back-half already

works) ? **100% `web/` + i18n (en/kn/hi), ZERO engine code** ; ready sketch in `DRY\_RUN\_CUJ.md`
 S2 + the workflow synthesis
 (`web/src/{api/types,api/client,hooks/useIngestJob,components/IngestPanel}.tsx` + `App.tsx`
 `videoId?state` + `ingest.\*` keys). Then opt-3 droplet (deploy fixed engine + GCS creds + #281 fix
 ? re-run CUJ on Linux). **? Repos SYNCED TO HEAD at session end:** khelsutra master `d0461fd` ;
 engine working tree `06ff4b7` ? NOTE the engine main checkout couldn't switch to `master`
 (sibling worktree `tier1-windowing` holds it), so it's the merged-but-remote-gone
 `feat/serving-m4-device-context` branch **reset --hard to origin/master + upstream repointed to
 origin/master** (safe: tracked tree clean + the 2 divergent commits were the M4 #283 squash,
 byte-identical to master; old tip tagged `\_pre\_sync\_m4\_44c9f01` + reflog). So `git pull --ff-
 only` works there now; the branch NAME is cosmetic (points at master head).

19 - **? B-P2 BUILT (in-UI bucket ingest) ? khelsutra

[#56](https://github.com/avidullu/khelsutra/pull/56) OPEN, CI GREEN (web+site+CF build),
 awaiting owner merge** (feature PR ? left for owner; precedent exists for me merging SPA PRs on
 owner's "merge as you go"). 100% `web/` + i18n, ZERO engine code (rides B-P1). New `web/src/{api/
 /types,api/client(processVideo/getJob),lib/errors(parseDetail+classifyError),hooks/useIngestJob(
 id?submitting?polling?done|error; recursive-setTimeout poll backoff+10min-ceiling+1 cid;
 success-shaped-failure?emptyResult),components/IngestPanel}` + `App.tsx` `videoId`
 useMemo`useState` + IngestPanel-when-!videoId + history.replaceState; `ingest.\*` en/kn/hi (AI-
 draft, native-review-pending). **123 web tests (+19), tsc/build green, browser-verified** (panel
 renders en+hi; live submit?friendly localized "is the engine running?"; cid logs). Live-fixed
 UX: bare 5xx/proxy-down ? friendly `networkError`, not "HTTP 500". `DRY\_RUN\_CUJ` S6 `box3` = ?
 (flips on a live **e2e-from-UI** vs a real engine). **? NEXT:** merge #56 ? live e2e-from-UI
 (local or droplet opt-3) ? flip box3; then localization DoD tail (B3 site `t()` shim; native
 kn/hi review).

20

21 **Checkpoint:** 2026-06-21 (NEW SESSION ? UX+l10n focus STARTED) ? ramped up via a
 research+verification **workflow** (web/+site/ string maps, Indic-l10n prior art, **5
 adversarial verdicts:** both surfaces **ZERO i18n today** ; recording-edu **already shipped** on
 the homepage gist + `/capture` ; `I18N\_PLAN` targets the **dead** `backend/api/static` so its
 P1/P3 wiring retargets to site/+web/). **? Isolated the work into ONE umbrella tracked project ?
 `khelsutra/docs/FRIENDLY\_AND\_MULTILINGUAL.md`** (A: non-tech UX + B: kn/hi l10n; S7 P-tracker =
 source of truth; **archive when en+kn+hi render on BOTH surfaces**). Pointers wired: khelsutra
 `docs/README` + `NEXT\_STEPS` ; engine `I18N\_PLAN` + `UI\_PROPOSAL` banners + `DOC\_STATUS` S3a/S3b rows
 (engine edits in an **isolated worktree off origin/master** ? engine checkout is a concurrent
 session's branch; their findings-#1/#2 PR **#274 merged mid-session** ? master `03fa951`).
 **Owner (AskUserQuestion) locked D6 (en+kn?hi) · D7 (OFL Noto self-host) · D8 (SPA-first) · D9
 (catalogs in khelsutra); only D-BACKEND-DEPTH deferred. 3 PRs OPEN:** khelsutra docs

[#41](https://github.com/avidullu/khelsutra/pull/41); engine pointers

[#275](https://github.com/Khelsutra/badminton-highlight-indexer/pull/275); **B1 SPA i18n
 scaffold + language switcher (en-only) ? [#42](https://github.com/avidullu/khelsutra/pull/42) ?
 CI GREEN (web+site pass), verified LIVE in-browser** (switcher EN|???? persists; content falls
 back to en with no kn catalog; active-highlight follows the *chosen* locale ? a bug caught by
 browser verification + fixed). B1 = react-i18next + ICU + detector + `en` catalog + the SPA's
 FIRST jsdom render tests (79 tests green); deps MIT/BSD-clean. Engine pointer PR built in an
 isolated worktree (now removed); engine checkout kept flipping branches (concurrent session
 active) ? never disturbed. **? NEXT slices (tracker S7): B2** (full ~50-string extraction + key-
 parity/no-bare-strings CI) ? **B3** (site `t()` shim) ? **B4/B5** (Noto fonts + kn catalog via
 Gemini-draft?review). **Also (owner ask, 2026-06-21): the en/kn/hi academy flyers already
 existed** (a sibling merged PR **#40** ? master `37053b9` ; flyer.html / flyer-kn.html / flyer-
 hi.html served at /flyer{,-kn,-hi}) but were UNLINKED ? **linked all three from the homepage
 footer + added VM-host clean routes ? [#43](https://github.com/avidullu/khelsutra/pull/43)
 (MERGED ? master `f121c7f`).** Then **owner asked for MORE flyer placement ?

[#44](https://github.com/avidullu/khelsutra/pull/44): "Flyer" in the desktop top-nav + the
 3-language links on own-model's foot** (44 site tests green; verified live). (Flyers use Google-
 Fonts Noto ? separate from the app's D7 self-host; not refactored.) **Owner #41 comment ? added
 a rigorous S6.5 "automatic language-negotiation" testing strategy** (matrix N1?N9: a kn/hi
 BROWSER auto-gets the respective version; DoD gate) **+ concrete browser-language detection
 tests in B1 [#42]** (N1/N3?N7; 84 web tests green). #41/#42/#275/#44 all merged. **? THEN
 owner: "take it to completion until you can run a screencast with a kannada-default browser ?
 right page renders ? switch hi ? en; merge PRs as you go." ? ? DONE for the `web/` SPA.** Merged
 in sequence (I merged each on green CI, owner-authorized): **B2 [#45]** (full ~50-string
 extraction + key-parity/no-bare-strings CI) · **B4 [#46]** (self-host OFL Noto
 Kannada/Devanagari, @fontsource-variable) · **B5 [#47]** (Kannada catalog) · **B6 [#48]** (Hindi
 catalog ? the extensibility proof: just a catalog + supported-list entry). **SPA now fully

en+kn+hi; 99 web tests green** (incl. N1/N2 browser-detection, B5/B6 render, key-parity kn==hi==en, ``<html lang>`` sync). **? SCREENCAST PRODUCED** (Playwright `locale: kn-IN` ? Kannada auto-renders ? switch hi ? en; `C:\Users\avidu\Projects\khelsutra-screencast\khelsutra-i18n-screencast.mp4` (9s) + per-locale frames). Tracker milestone ? [#49]. ? **kn/hi are AI-DRAFTED (me) ? native-speaker review pending** before wide public launch. **? REMAINING for full \$9 DoD (NOT archive-ready):** **B3** (the `site/` marketing surface ? homepage/capture still English-only ? via the shared `t()` shim) . **A3** recording-education kn/hi . the **non-tech UX A-items** (A1/A2/A4/A6...) . native review. Playwright installed in isolated `C:\Users\avidu\Projects\khelsutra-screencast\` (kept out of web/ deps). **THEN owner polish batch (2026-06-21) ? all merged:** (1) **OFFLINE/self-contained verified + guarded ? [#50]:** the SPA is air-gap-renderable (Noto = local woff2 via @fontsource, `@font-face` url?local `/assets`, no Google-Fonts/CDN, `web/index.html` no external link); CI guard added so it can't regress. *(The marketing `site/` still uses Google-Fonts ? hosted/online surface, not the offline tool; self-host = optional follow-up.)* (2) **SCREENCAST uploaded ? `gs://khelsutra-rally-corporus/promo/khelsutra-i18n-screencast.mp4`** (+ local). (3) **Localized SAMPLE-QUERY PROMPT (en/kn/hi) ? [#51]:** `query.hint` = instructive "type `over 10s` for rallies >10s ? or tap a sample" + `query.try` labels pills "Sample queries"; verified live kn. (4) **\$6.5 testing rule:** i18n testing is now a MANDATORY merge-gate for any future query-parsing/semantic-understanding/NL work (parser is English-grammar-only today, A9). **103 web tests green.** NO open PRs. [[next-session-ux-localization]] [[doc-status-register]] [[working-agreement-tracked-project-docs]]

22

23 **Checkpoint:** 2026-06-21 (DRY-RUN CUJ + a CRITICAL ENGINE FIX) ? owner: "stitch the UI to the tool; dry-run a **GCS-bucket video** via the UI ? **reel downloaded local**; repeat in Hindi. Run locally (opt 1) then droplet (opt 3); report both; identify fixes + merge." Chose **A-now-then-B**, **Gemini**. **Tracked goal ? `khelsutra/docs/DRY\_RUN\_CUJ.md`** ([#53]/[#54]). **? PATH-A LOCAL RUN DONE (en + hi):** isolated engine off origin/master (fresh DB/output worktree) + `.env` Gemini key ? `gsutil cp gs://khelsutra-rally-corporus/videos/mahadevpura\_smoke\_180s.mp4` ? `POST /api/process {gemini}` ? **12 rallies** (`video\_id 05e0e577?`) ? SPA `npm run dev` (Vite proxies /api?:8000) `?video=05e0e577?` ? loaded 12 real rallies ? **Build ? /api/compile ? reel (1080p, 76.5s) DOWNLOADED to `~/Downloads/khelsutra-dryrun/highlights\_local\_{en,hi}.mp4`** (50.9MB). **Hindi repeat ?** (page/cards/"??? ?"/"??? ?" all Hindi; reel downloaded). **?? CRITICAL BUG FOUND + FIXED ? engine [#281](https://github.com/Khelsutra/badminton-highlight-indexer/pull/281) MERGED:** running the documented `python -m backend.main --server` runs main.py as `\_\_main\_\_` (sets `\_\_main\_\_.db\_instance`), but `job\_worker`/`uploads` read `import backend.main` ? a SEPARATE module whose `db\_instance` stays None ? the async-job drain crashes silently on `None.claim\_due\_job()` ? **every /api/process job hangs in `queued` forever**. Latent: unit tests monkeypatch the executor inline; the UI was verified vs a STUB engine. Fix = dispatch `\_\_main\_\_`?`backend.main`; + `tests/test\_server\_entrypoint.py` (real subprocess guard, FAILS-without/PASSES-with). Full engine suite green (2m23s). **Unblocked the run via the single-module invocation** (`python -c "...; from backend import main as m; m.main()"`). ? **OPT-3 DROPLET (168.144.126.176) NOT YET RUN:** reachable (Linux, py3.12, ffmpeg, 2-core/3.8GB) but **un-provisioned** (no `~/rally` engine checkout) ? needs a from-scratch engine install OR a proper deploy of the FIXED engine (the fix applies there too). **? NEXT:** opt-3 droplet run (deploy fixed engine ? re-run CUJ); Path B (in-UI bucket ingest). ? **CORRECTION (next-session verify): the "/api/process is local-path-only / net-new gcs ingest" note here was STALE** ? `api/process` already accepts `gs://` (M2/#280, tested); Path B is **SPA-only (B-P2)**, no engine endpoint. (The "worker writes CSV to the bucket not the API DB" bit is the *separate* `worker infer` CLI, not the API path ? that's what the stale claim conflated.) Docs fixed ? khelsutra [#55]; see the NEXT SESSION FOCUS block above + [[lesson-verify-engine-surface-before-scoping]]. [[monetization-plan]]

24

25 **? HANDOFF (2026-06-20) ? KHELSUTRA.GURU brand site: shipped, live, portable. Resume the product threads below.** Separate track from the indexer rally-detection handoff that follows. Full durable record + ramp-up: [[khelsutra-domain-launch]]. Repo = **`avidullu/khelsutra`** (NOT this indexer repo).

26 - **? START HERE (2026-06-20, latest ? big session, ALL MERGED, NO open PRs).** Per-rally **MVP/M0 COMPLETE** + a full product/promo wave. **khelsutra master `a7214ad`** . **engine (Khelsutra/badminton-highlight-indexer) master has M1 [#268]**.

27 - **Per-rally MVP (khelsutra `web/` SPA):** P3 `parseQuery()` query bar [#28] (length/order/count/shot_count; greys shot-type/player/score "not available yet" ? never faked; **13 adversarial-review findings folded in**) . P4 `SelectionFooter`+honest `min\_shot\_count` warning [#30] (`web/src/lib/floor.ts` mirrors `main.py:362-383`; no-shot_count exempt; all-below ? Build DISABLED) . P5 real load `?video=<id>` + Build?`compileReel`?`<video>` preview+download

```

[#31]. Full obs (`query.*`/`build.*`/one cid). **Then: lead-with-LENGTH, shot-count flagged
EXPERIMENTAL** [#33] (cards show `m:ss long`; amber "experimental" chip; duration-led examples;
badge `P5`?`alpha`; parseQuery still accepts shot-count, de-featured). Suite **75 green**.
28 - **Engine M1 [#268] (GCP/org repo, built in an ISOLATED git worktree off
origin/master so the concurrent session's checkout was untouched):** `GET /api/health`
({status,sport,default_segmenter}) + `POST /api/compile` now returns **`download_url`**
(additive; `filepath` unchanged). Full engine suite **1019 passed/7 skipped**, ruff+mypy clean,
CI green (coverage floor incl. torch). *(SPA can adopt `download_url`/offline-banner as a
follow-up.)*
29 - **GCS round-trip (reuse artifacts, CPU/no-GPU):** pulled
`mahadevpura_smoke_180s.mp4` + intervals from the bucket ? **stitched a curated top-10-longest
highlight** (ffmpeg) ? **put it back** at `gs://khelsutra-rally-
corpus/highlights/curated_top10_longest.mp4` (61s).
30 - **3 screencasts (Playwright?webm?ffmpeg):** `promo_cuj.mp4` (old UI, real footage) .
**`full_cuj_real.mp4`** (NEW duration-first UI + REAL curated reel ? the definitive full-CUJ) ?
both in `gs://khelsutra-rally-corpus/promo/` (real footage ? bucket, NOT public repo) .
**`site/public/video/tool-cuj.mp4`** (NEW UI, **footage-free** branded-placeholder reel ? safe
to self-host). Local copies in `badminton-highlight-indexer/artifacts/promo/` (gitignored).
31 - **Site (excite-to-join) [#34/#35]**: **after sign-up the reveal now shows ***(1) the
tool in action** (self-hosted footage-free `tool-cuj.mp4`) **+ (2) a ~30s reel preview**
(`REEL_SECS` 90?30, gated nocookie YouTube ? real faces stay in the unlisted reel, not the
repo); `revealAfterSignup`; coherent CTA/teaser copy; `server.js` serves `.mp4` as `video/mp4`;
`.mp4`/`.gif` marked binary. Site suite **34 green**; **droplet re-synced + parity GREEN**
(index+mp4 byte-identical CF?VM:8787).
32 - **GPU note:** could NOT spin a fresh GPU VM ? `GPUS_ALL_REGIONS=1` slot held by
the concurrent session's `rally-gpu` M2 training (did NOT disturb). Validated CUJ against the
first-L4-run REAL artifacts instead. A fresh live `/api/process` GPU inference stays
**deferred** until the GPU frees (low-risk existing path). **Owner observation ? [[machine-
pooling-idea]]** (pool GPUs vs spin/kill per task).
33 - **Owner decision flags (PUBLIC promo):** `full_cuj_real.mp4`/`promo_cuj.mp4` show
real court footage (faces) ? for public use beyond the gated site, confirm. The footage-free
`tool-cuj.mp4` is safe.
34 - **COLD-START RUN (2026-06-21, owner-approved; GPU freed ? the other session
deleted `rally-gpu`):** provisioned a FRESH **L4** `rally-coldstart` (T4 stocked-out asia-
south1-a ? L4 -c), built engine+native-WASB env in **~7 min** (`bootstrap
--gpu`+`gpu_setup.sh`+weights; main torch 2.6+cu124 & wasb-env 1.11+cu113 both cuda True),
`worker infer` GCS-in?GCS-out, **VM DELETED ($0, ~$0.40 total)**. **Cold-start MECHANISM fully
validated end-to-end.** **BLOCKER: native WASB returned 0 rallies** on mahadevpura-1 (full-
res), GX030094_proxy, AND the known-good **smoke clip (gave 22 in the first L4 run)** ? a
**fresh-build CONFIG/wiring issue, NOT the CUJ**: my minimal reconstructed `config.json`
(`indexing.detector_impl=native`, `default_segmenter=wasb_hybrid`, no extra fusion/gating
tuning) doesn't reproduce the engine session's TUNED `worker infer detector_impl=native`
invocation. **THE missing piece for a repeatable cold-start = the engine session's EXACT
native-WASB config/cmd** (the one that gave 22 rallies). Side finding: `mahadevpura-1.MP4` is
full 1080p/~25Mbps; the working clips are downscaled `*_proxy` ? feed proxies (or add a
downscale step). ? The diagnostic infer **clobbered**
`outputs/05e0e577?/{intervals.csv,report.json}` ? **RESTORED to 22 segments** from local
`artifacts/promo/intervals.csv` (reel untouched; junk 0-seg dirs removed). **ROOT-CAUSE OF "no
logs" FIXED ? engine [#271](https://github.com/Khelsutra/badminton-highlight-indexer/pull/271)
MERGED** (isolated worktree, full gate 1027 green): `backend/worker/__main__.py` **never
configured logging** ? all pipeline diagnostics (detector choice, "Running native WASB
inference", availability failures) were silently dropped at default WARNING. Now:
`_setup_logging()` at the CLI entry (INFO; `force=True` so it owns the level) + a `-v/--verbose`
DEBUG flag (native cmd, frame counts); and `cmd_infer` emits a **LOUD WARNING on a 0-segment
success** (video_id + segmenter + detector_impl + likely causes + "re-run -v"). So **next cold-
start, a 0-rally run is fully analysable from the log** ? that's the diagnostic that was missing
this time.
35 - **VIDEO?HIGHLIGHTS REVIEW TABLE (owner asked):** in `gs://khelsutra-rally-corpus/`
**only `mahadevpura_smoke_180s.mp4` is fed** (id 05e0e577; **22 rallies, 94.2s play, ~48% dead-
time cut**) ? reels `outputs/05e0e577/highlight_smoke_l4.mp4` (full ~95s) +
`highlights/curated_top10_longest.mp4` (top-10, 61s). **mahadevpura-1.MP4` (GT 10) +
`GX030094_proxy.mp4` (GT 41) staged, NOT yet fed** (cold-start blocked above). Local review
copies: `badminton-highlight-indexer/artifacts/review/` (source + both reels). 7 GT
`.rallies.csv` in `labels/`.
36 - **OPEN owner-gated (unchanged):** appliance $10 decisions; move khelsutra site repo

```

? Khelsutra org (needs Cloudflare GitHub-App reconnect). ****Next options:**** (1) ****get the engine session's exact tuned native-WASB `worker infer` config/cmd**** ? a cold-start that actually yields rallies+reel (the one missing piece); (2) SPA adopt `download\_url`+offline-banner (uses M1); (3) engine P7 (`/api/videos/{id}/rallies` alias); (4) polished public promo cut. **[[machine-pooling-idea]]**

37 - ****? SESSION WRAP (2026-06-20 late) ? site fully LIVE on Cloudflare + the DO droplet twin (`/srv/khelsutra-site`), byte-in-sync (parity green); PRs #1?#26 merged, NO open PRs, master `9f0942b`. ** This session shipped: ****per-rally MVP P1+P2**** (`web/` SPA: Vite/React/TS/Tailwind + typed engine client + `RallyGrid` + `useRallySelection`; 16 vitest; ****P3 now done ? see START HERE above****) . ****3 tracked design docs merged**** (`PER\_RALLY\_SELECTION.md` IN-PROGRESS, `LOCAL\_APPLIANCE\_ARCHITECTURE.md`, `ANNOTATE\_AND\_TRAIN.md`) . live ****~/own-model**** page (real annotation-tool screenshots [no people] in `site/public/img/`, sample reel ****gated**** behind request-access, "default works out of the box / BYO-training ****optional****" framing, ****? KhelSutra home**** back-link) . ****provider-portable**** `site/scripts/provision-vm.sh`+`teardown-vm.sh` (LIVE-tested: fresh DO VM ? operational ? state-export ? destroy) + the ****git -c core.autocrlf=false archive**** CRLF fix . ****homepage sample reel**** gated post-signup (random-start `youtube-nocookie` + 90s auto-close) . ****rally-annotator` TRANSFERRED ? `Khelsutra` org\*\*** (links repointed). **\*\*Site CI = 32 green\*\*** (`npm test`) + scheduled live-parity workflow. ****GPU deferred to the owner's Windows+WSL box**** (DO has no GPU SKUs + GCP quota=0; no Gemini e2e per owner). ****OPEN owner threads:**** per-rally ****P3+****; appliance ****\$10 decisions**** (edge-auth / retention / per-video cost cap / GPU host / P5 bucket-worker bridge); ****move sister repos ? `Khelsutra` org as PRIVATE**** (verified feasible ? free org allows 10000 private repos + collaborators, owner stays admin; ENGINE-repo move needs coordinating w/ active sessions ? see `khelsutra/NEXT\_STEPS.md` "Org & repos"); optional VLC-CSV annotate walkthrough screenshots. ? Droplet STILL a shared testbed (only `/srv/khelsutra-site` touched; never the engine `~/rally`/MinIO/GCS). ****Ramp-up:**** this block ? **[[khelsutra-domain-launch]]** ? `khelsutra/NEXT\_STEPS.md` ? the 4 design docs (`PER\_RALLY\_SELECTION` / `LOCAL\_APPLIANCE\_ARCHITECTURE` / `ANNOTATE\_AND\_TRAIN` / `PORTABILITY`). Droplet SSH: `~/ssh/khelsutra\_droplet` ? `root@168.144.126.176`; sync recipe in `PORTABILITY.md`.**

38 - ****? ORG CONSOLIDATION DONE (2026-06-20 late):**** transferred ****`badminton-highlight-indexer` (engine), `sports-data-collector`, `sports-obsreport` ? the `Khelsutra` GitHub org**** (all stay ****private****; verified `Khelsutra`/repo private). Safe: no open PRs, no forks, already-private (no visibility change); GitHub ****redirects git ops**** so active quality-iteration sessions' push/pull/CI keep working (local files untouched). ****Local remotes repointed**** to `https://github.com/Khelsutra/<repo>.git` (engine ls-remote verified). ****rally-annotator` stays PUBLIC\*\*** (OSS MIT plugin the site links to ? do NOT make private). **\*\*`khelsutra` (site repo) NOT moved\*\*** ? a transfer breaks the **\*\*Cloudflare Workers-Builds git connection\*\*** (auto-deploy) until the CF GitHub App is installed on the `Khelsutra` org + the Worker reconnected (owner dashboard step) ? **\*\*pending owner\*\***. Engine harness/memory path UNCHANGED (keyed to the local dir name `badminton-highlight-indexer`, not the GitHub URL). ? Could not read org Actions policy (gh token lacks `admin:org`); engine CI history intact at the new path (last run = success) + free orgs default Actions-on ? verify CI on the next engine PR; enable in org settings if needed.

39 - ****DONE + LIVE this session (PRs #1?#8, all merged, repo clean):**** owner bought khelsutra.guru ? built & launched end-to-end. On ****Cloudflare**** (Worker `hidden-king-khelsutra` + D1 `khelsutra-access`; ****git-connected ? merge `master` auto-deploys****; CI runs `npm test` = 20 green): landing page + `/capture` checklist + OG image + a ****segmented Request-Access form**** (academy/individual + GPU + storage Qs + free text) ? persists to D1 ****and emails a summary to avi.dullu@gmail.com per registration**** (Email Routing send binding ? VERIFIED live in inbox). Valid SSL; `hello@` Email Routing verified.

40 - ****PORTABLE + proven****: identical form runs on a dependency-free ****Node + SQLite**** host (`site/server.js`), ****live + browsable on the owner's DO droplet ? http://168.144.126.176:8787**** (systemd `khelsutra-site`, isolated `/srv/khelsutra-site`; SSH key `~/ssh/khelsutra\_droplet`). Both-hosts compat ENFORCED in CI (`site/src/server.test.js` boots the real server). ? Droplet is a SHARED storage testbed ? another session runs MinIO 9000/9001 + GCS emu 9023 + `/root/do\_gcs\_`; STAY CLEAR. Docs: `docs/REQUEST\_ACCESS\_FORM.md`, `docs/PORTABILITY.md`, `NEXT\_STEPS.md`.

41 - ****? PENDING / next threads (owner ideas this session):**** (1) ****Marketing video**** ? use Grok/Gemini image-gen ? a short video themed for ****serious enthusiasts**** showing the whole record?reel process ? embed on the site. (2) ****Interactive per-rally selection**** (big product feature) ? the engine detects rallies ? present them in a UI so the user ****picks which rallies go into the highlight**** (dynamic, NOT a fixed setting). ? ****Design doc MERGED ? [khelsutra#11](https://github.com/avidullu/khelsutra/pull/11); owner APPROVED ? decisions D8?D11 locked (status IN PROGRESS). ? CI regression-guards MERGED ? [khelsutra#12](https://github.com/avidullu/khelsutra/pull/12): `site/src/deploy-config.test.js` (pins wrangler.toml deploy invariants ? the #5 name-drift class), `site/src/assets.test.js`**

(static-site content contract), `.github/workflows/parity.yml` (scheduled+dispatch live-parity, NON-blocking); `npm test`?`node --test` discovery; suite 20?28 green; live-parity dispatched ? BOTH hosts byte-in-sync. ? **P1 MERGED ?

[khelsutra#14](https://github.com/avidullu/khelsutra/pull/14)** (squash `aac7e0c`): `web/` SPA scaffold (Vite6+React19+TS+Tailwind4) + **typed engine API client**

(`listRallies`/`compileReel`/`highlightUrl`/`getConfig`+`minShotCount`, grounded+verified; honest types flag `server/receiver/events` as motion fabrications) + **correlation-id structured logging** (`web/src/lib/log.ts`, observability seed, loud failures) + **8 vitest green** + new **`web` CI job** (npm ci?typecheck?test?build, gated lockfile); live-verified renders in-browser with ZERO console errors. Isolated (no README/NEXT_STEPS edits ? no #13 conflict). ? **P2 DONE ?

[khelsutra#21](https://github.com/avidullu/khelsutra/pull/21):** `RallyGrid` (honest fields only) + `useRallySelection` hook over pure `lib/selection.ts`+`lib/format.ts` (default ALL-on-per D8); 16 vitest green; live-verified (deselect ? footer recompute 6/1:22?5/0:53). **NEXT = P3** `parseQuery()` ? P4 footer+`min\_shot\_count` warning ? P5 build?compile?download. All on existing endpoints.** Built via a multi-agent design Workflow (research real engine contract ? 3 approaches ? synthesize) with **every load-bearing claim personally re-verified vs engine source**. Key finding: **MVP needs ZERO new engine endpoints** ? rides existing `POST /api/query` ? select `segment\_ids` ? `POST /api/compile` ? `GET /api/highlights`. Recommended hybrid: selectable rally **grid** (workhorse) + **deterministic `parseQuery()`** (no LLM; length/count/order/shot_count ONLY) writing into ONE shared selection set + sticky compile footer. Honest scope: shot/player/score = roadmap (no column/detector);

`server`/`receiver`/`events` = `motion`-only fabrications ? hidden; default `gemini` writes real-ish `shot\_count`; per-rally `preview` blocked until a source/clip stream endpoint lands (deferred milestone = biggest UX-vs-effort lift). Owner decisions RESOLVED (D8?D11, 2026-06-20): (a) default selection **ON** (curate-by-removal); (b) preview-before-select **NOT** a ship-blocker (text-metadata MVP; preview = P9); (c) **job-tethered** MVP (no `GET /api/videos`); (d) `min\_shot\_count` = **warn, don't override** (override deferred ? Launch Report). Doc = `khelsutra/docs/PER\_RALLY\_SELECTION.md`. Pairs with the "ask for a reel" conversational builder already teased on the site. (3) productionize droplet (TLS + `vm.khelsutra.guru`), VM-path SMTP email (currently noop), Always-Use-HTTPS/HSTS, Web Analytics token, `/capture.html`?`capture` footer link, delete any orphan `khelsutra-site` Worker, clean CF D1 test rows (`deploy-test@`/`email-test@`). Revoke the droplet SSH key when done. [[paper-coauthor-aim]]

[[monetization-plan]]

42 - **? APPLIANCE DESIGN ? doc SHIPPED for review ?

[khelsutra#13](https://github.com/avidullu/khelsutra/pull/13) (DRAFT, owner-gated; CI green) ? "locally-installed, UI-powered rally indexer + reel generator" (KheISutra Studio).** Owner target (HEADLINE CUJ): a **hosted website** + a **NEW GPU-enabled machine (provider-agnostic: DO GPU droplet / GCE GPU VM / Cloud Run+GPU)** running the decoupled e2e; from the UI the user picks an input ? **Google Drive / bucket URI (S3/GCS/DO Spaces) / local upload?bucket** ? then **generates + watches + downloads** highlights **completely from the UI**. Framed as a **deployment PROFILE of the decoupled-compute arch** (engine `backend/storage/` StorageBackend + `python -m backend.worker {infer|train}` bucket pipeline; PRs #246?#254 on engine master, 936 green; **the CPU droplet 168.144.126.176 already hosts BOTH the site `/srv/khelsutra-site` AND the engine `~/rally`**, GPU MOCKED via stub). ? **OWNER PRIORITY (2026-06-20): MOVE TO A GPU-ENABLED MACHINE ASAP ? provider-agnostic: DO GPU droplet / GCE GPU VM / Cloud Run+GPU** (storage stays the S3/GCS bucket; GCS = storage, GCE/Cloud Run = compute ? do NOT conflate). Critical-path deps that are NOT mine to do unilaterally: (owner) stand up the GPU compute target + provide the cloud API token/creds ? DO runbook = engine `docs/SETUP\_NEW\_MACHINE.md` \$2/\$9 + `deploy/digitalocean/`; **Cloud Run / GCE need the `backend.worker` CONTAINERIZED (likely a GAP; WASB's torch-1.11 conda env makes the image non-trivial) ? verify**. Cloud Run+GPU (scale-to-zero) is the strong fit for the per-video-billing economics (engine DECOUPLED_COMPUTE #246). (engine session) build **`NativeWasbRunner` (M3.5)** for real (non-mock) WASB. The **UI/appliance work runs in PARALLEL** (rides the existing FastAPI: chunked upload, `/api/process`?`/api/jobs`, `/api/query`, `/api/compile`, `/api/highlights`). ? Engine API has **NO auth (CORS `\*`)** ? **EDGE AUTH required before any public exposure**. ? Google **Drive ingest may be a `gdrive` STUB** in `backend/storage/` (#249) ? verify before promising. Doc ? `khelsutra/docs/LOCAL\_APPLIANCE\_ARCHITECTURE.md` (PR pending). The merged per-rally selection (#11) = ONE screen of this operator console. CUJ/live-test plan is a first-class doc section (owner asked re: live CUJ replay). **RECOMMENDATION (B+C hybrid):** reverse-proxied **UNIFIED ORIGIN** ? `web/` SPA behind edge auth, `/api/\*` ? engine `127.0.0.1:8000` (CORS `\*` moot; gate = tenancy boundary; bind verified `main.py:638`) ? driving the **EXISTING** FastAPI loop (upload?process?jobs?query?compile?highlights) with `storage=local`; **compute-as-a-setting**, bucket/remote-GPU worker = reserved scale path. **Verified-honesty refinements:** native WASB PARTIALLY wired (`fusion\_hybrid` builds `NativeWasbRunner` :86; `trajectory\_hybrid` refuses :57-63) + owner-byte-parity-gated; `ai\_max\_workers` default=**5** not 1; chunked-upload

endpoints EXIST (`uploads.py:64/89/123`) but `app.js` never calls them (client-wiring only).
NET-NEW engine endpoints to COORDINATE with the decoupled sessions: `/api/capabilities`,
`/api/videos`, `/api/reels`, `/api/health`. **\$10 owner decisions PENDING:** edge-auth (Caddy
basicauth vs cloudflared+Access), retention defaults, per-video cost ceiling, GPU host choice,
P5 bridge ownership. **Observability = first-class \$5** (correlation-id e2e, structured logs per
path, loud failures, health/metrics) per [[feedback-launch-quality-bar]]. Doc =
`khelsutra/docs/LOCAL\_APPLIANCE\_ARCHITECTURE.md`. **DELIVERY LOCKED (D10, 2026-06-20):** self-
hosted **webserver + UI the user spins up** (NOT a desktop app; supersedes
`DESKTOP\_APP\_PLAN.md`); user points at local storage+GPU/CPU OR cloud, UI drives the tool.
PROVISIONING ACCESS (2026-06-20): `doctl` **INSTALLED v1.162.0** (winget
`DigitalOcean.Doctl`; at `%LOCALAPPDATA%\Microsoft\WinGet\Links\doctl.exe`, on PATH) ? **doctl
AUTHED** (team "Khelsutra Guru", droplet-limit 3; 1 used = `claude-do-rally-test` blr1 CPU
@168.144.126.176). ? **GPU BLOCKED ON BOTH PROVIDERS (verified 2026-06-20):** DO catalog has
NO GPU sizes (17 CPU-only sizes; DO GPU is access-gated + region-limited to NYC/TOR/ATL, not
blr1) AND **GCP `GPUS\_ALL\_REGIONS` quota = 0** (Compute API enabled). Both need an **owner-
initiated, provider-approved access/quota request** ? NOT CLI-doable (GCP quota-increase =
fastest realistic path since gcloud is authenticated as avi.dullu@gmail.com, project `gen-lang-
client-0306026826`). ? \*\*DO NOT DELETE the old droplet\*\* (`claude-do-rally-test`) ? it's the
SHARED decoupled-session testbed (`~/rally` engine, MinIO 9000/9001, GCS emu 9023, validated DO
Spaces e2e); deleting destroys their env. Site-migration + droplet-deletion is MOOT until a GPU
VM exists anyway. ? **Real product e2e (upload?index?reel) is testable NOW via Gemini (no GPU
needed)** ? GPU is only for the owned/native-WASB model + the train-your-own-model USP (also
pending the engine session's `NativeWasRunner` M3.5). ? **Annotate+train USP design SHIPPED ?
[khelsutra#15](https://github.com/avidullu/khelsutra/pull/15)** (DRAFT, owner-gated; CI green).
Plan: ONE marketing page `site/public/own-model.html` (Annotate/Train/Setup sections, clones
`capture.html`, links MIT rally-annotator) + a design-tracked owner-gated SPA correction screen
(NO "Train now" button ? training is offline `backend/worker train` CLI). VERIFIED: annotator
CSV (`rally\_number, start\_time, end\_time, ending\_reason, sport, shots\_count`) ingests with **NO
transform** (synth `DictReader row["start\_time"]`, `stub\_runner.py:139` + `gt\_loader`); ONE
prereq = **MD5+stem video?label join** (never re-encode after annotating). Honest: annotator
real+shipping (v1.6.4 MIT VLC plugin); train half = open CLI + roadmap (GPU-free=CPU **mock**
only; real needs GPU + in-flight native WASB; ship-gate `gate\_verdict` `harness.py:98` refuses
commercial ship; `label\_source=human` + `no\_paid\_llm\_training\_labels` attested :178/:189-192);
annotator repo has **ZERO images ? all screenshots TO-CAPTURE** (consented, no minors, NEVER
fabricate); USP = ownership/licensing/privacy/consent/minors-excluded, **NOT accuracy** (no
benchmark vs paid LLM). **GPU DECISION (owner 2026-06-20):** use the **Windows+WSL GPU box** for
real testing; **DEFER the cloud GPU VM** ; **NO Gemini e2e now** ; **NO droplet
migration/deletion** . Real WASB e2e ? owner's box when run deliberately. **#13 (appliance) + #15
(annotate/train) MERGED** (owner-approved). ? **PROVIDER-PORTABILITY now first-class + LIVE-
TESTED ? [khelsutra#16](https://github.com/avidullu/khelsutra/pull/16) MERGED:**
`site/scripts/provision-vm.sh` (fresh Ubuntu VM on ANY provider ? `git archive` push [private-
repo-safe, no secret on VM] ? Node24 + systemd + firewall + verify) + `teardown-vm.sh` (EXPORT
SQLite state to local backup BEFORE destroy ? clean resume via `RESTORE\_DB`) + PORTABILITY
runbook + CI `bash -n` parse-check. **PROVEN on a throwaway DO droplet (then destroyed):** bare
? operational ? public HTTP 200 ? state-export ? re-provision-with-RESTORE_DB preserved the row
(`rows:2`) ? droplet+temp-key removed. So destroy/migrate-across-providers is now safe +
reproducible (owner principle: "deleting is fine IF we can resume cleanly via scripts"). ? Still
DO NOT delete the shared `claude-do-rally-test` droplet without coordinating the decoupled
session (their `~/rally` engine env reproduces via the ENGINE repo's `SETUP\_NEW\_MACHINE.md`, not
mine). ? **UI PROPOSAL rendered** (visualize mockup): an ***"Own your model"*** section for the
EXISTING landing page ? Annotate-in-VLC (live) / Train-your-own (in-build, honest) /
Yours+private ? links the MIT annotator + a `own-model` setup guide; ? **own-model page LIVE ?
[khelsutra#22](https://github.com/avidullu/khelsutra/pull/22)** (`site/public/own-model.html`:
Annotate=live / Train=in-build / Setup; honest framing, GitHub-linked, **placeholder**
screenshots; `own-model` alias + nav/footer cross-links; deployed to BOTH hosts, parity green;
real screenshots = owner-gated **P5**). ? **own-model reel REUSED as click-to-play ?
[khelsutra#24](https://github.com/avidullu/khelsutra/pull/24):** owner asked to reuse the
existing reel rather than capture screenshots ? the video is a REEL (engine output), NOT a shot
of the annotator UI, so it can't replace a tool screenshot, but it's embedded as a **branded
click-to-play** (ink poster + play button, NO footage/faces until tap ? `youtube-nocookie`
IFrame-API player, **random start ?90s before end**). Footage hidden until tap (privacy-
consistent w/ homepage); verified live both hosts (`facade=4`, static-iframe=0, thumbnail=0,
parity green md5 `90de6a9e`). ? **SUPERSEDED by
[khelsutra#25](https://github.com/avidullu/khelsutra/pull/25) ? more coherent layout (owner
call):** the reel (shows players) is now GATED in ONE place only (homepage request-access); the

own-model page no longer embeds it. Instead own-model shows the ****real annotation-TOOL screenshots**** (owner provided 2, NO people: ``site/public/img/annotate-enable.png`` = VLC View? Rally Annotator, ``annotate-tool.png`` = the Mark START/END/Save dialog) + a one-line ****"request access to see a sample output" pointer + a "you don't have to train anything ? a ready-made default works out of the box; BYO-training is OPTIONAL" note** (keeps non-tech users engaged). Tests assert screenshots-present + no-reel-embed + default/optional framing. Verified live BOTH hosts: images HTTP 200 image/png, md5 ``b6350339`` identical local/droplet/CF (binary integrity through git-archive), own-model ``iframe=0/youtube=0``, parity green. (Reel still lives gated on the homepage: random-start nocookie + 90s auto-close after signup.) ? ****rally-annotator ? Khelsutra org:**** owner created the ``Khelsutra`` GitHub org; transfer is FEASIBLE (owner = org admin, no name conflict) ? ? ****DONE: TRANSFERRED**** to ``Khelsutra/rally-annotator`` (org owns it; old ``avidullu/?`` URL auto-redirects). Site/doc links repointed ?

[khelsutra#23](https://github.com/avidullu/khelsutra/pull/23); deployed to BOTH hosts, ``/own-model`` parity green (md5 ``c0af70fb``). Repo ``git remote`` for any local rally-annotator checkout should be updated to the Khelsutra URL. GPU still deferred to the Windows+WSL box (no cloud GPU; no Gemini e2e). ? ****DEMO REEL ? SIGN-UP-GATED + on BOTH hosts (#17 open ? #18 gated ? #19 fix).**** The owner's UNLISTED reel ``ajTCZ7pJW_o`` is NOT shown openly; instead a "Sign up to see it in action" teaser + on a successful request-access submit it reveals a ``youtube-nocookie`` autoplay embed that ****AUTO-CLOSES in 90s**** (``#reel-reveal`` + live countdown; [khelsutra#18](https://github.com/avidullu/khelsutra/pull/18) merged; verified live submit?reveal on the Node host; ``assets.test.js`` asserts GATED-not-open; suite 29 green).

****Deployed to BOTH:**** Cloudflare auto-deployed; ****droplet `/srv/khelsutra-site` twin re-synced to master ? PARITY GREEN**** (``OK / 231dd90`` · ``OK /capture 2ada2cf``; access.db + self-contained Node 26 preserved; only ``/srv/khelsutra-site`` touched). ? ****CRLF-drift bug found live + fixed ?** [khelsutra#19](https://github.com/avidullu/khelsutra/pull/19): ****Windows `core.autocrlf=true` overrides `.gitattributes eol=lf` in `git archive` ? CRLF on the VM ? byte parity drift vs CF's LF; `provision-vm.sh` + `parity-check.sh` now pin `git -c core.autocrlf=false archive` (LF cross-platform; no-op on Linux). Privacy now stronger: the clip only plays after signup (not for every visitor/crawler); YT stays owner-controlled. ? ****#20 (random start + join FAQ):**** the reel now starts at a ****RANDOM point 90s before the end**** (YT IFrame Player API: ``getDuration()`` ? ``random [0,dur-90]`` ? ``seekTo``; falls back to 0 if dur unknown; verified live = 35s then 20s across two sign-ups) so people don't all see the same chunk; the ****"How do I join?" FAQ**** now points to the request-access form (``#join``, no ``mailto``) ? consistent with the actual signup. Both hosts re-synced LF-safe ? ****parity GREEN**** (``848312?``); site suite ****30 green****.**

[[feedback-launch-quality-bar]] ****`gcloud` IS authed**** here (avi.dullu@gmail.com, project ``gen-lang-client-0306026826`` ? likely needs GPU quota + billing verification before GCE/Cloud Run). Owner prefers ****DO first**** for GPU. Creating paid GPU infra = confirm-first every time; I propose exact cmd + hourly cost + teardown, owner OKs. [[monetization-plan]] [[khelsutra-domain-launch]]

43
44 ---
45
46 ****? HANDOFF (2026-06-20) ? DECOUPLED_COMPUTE: M3.5 native runner MERGED + FIRST REAL GPU RUN DONE on GCP L4 (M1? M3.5? M2?). VM deleted (clean slate). Resume/repro from `deploy/gcp/README.md` + `docs/DECOUPLED\_COMPUTE/DEPLOY\_REPORT\_GCP\_2026-06-20.md`.**** Resumes the engine-side dependency the appliance block above was waiting on.

47 - ? ****M3.5 MERGED ? [#256](https://github.com/avidullu/badminton-highlight-indexer/pull/256):**** ``NativeWasbRunner`` (``backend/pipeline/detectors/native_wasb_runner.py``) runs ``wasb_infer.py`` NATIVELY on a Linux GPU box ? no WSL/conda-activate/mnt; ``detector_impl="native"`` seam; ``device`` cuda/mps/cpu (cuda default = byte-parity with WSL); ``weights/repo/python/scratch env-first`` (``WASB_WEIGHTS`/`WASB_REPO_DIR`/`WASB_PYTHON`/`WASB_SCRATCH``). ``wasb_infer.py`` got ``--device`` + device-keyed cache (back-compat). ****`worker train --trajectory-source detector`**** does REAL per-video extraction (fps PROBED via ``ffprobe`` + skip-on-miss, video-ext filter, graceful native-misconfig). Suite ****998 green****, ruff+mypy clean, cov 83.7%. Adversarial multi-agent review run; fixes folded in. ****Parity ? VALIDATED 2026-06-20 on the owner's RTX 2070 (via WSL ? a Linux+CUDA proxy): native==WSL 1194/1195 frames, two native runs byte-identical (deterministic); gate = ~1e-3px TOLERANCE (cuDNN sub-pixel cross-GPU), NOT byte-exact. Native runner now DEFAULTS to streaming (#259). See the validation checkpoint below.****

48 - ? ****DO?GCP PIVOT (ADR-010) + runbook ? [#257](https://github.com/avidullu/badminton-highlight-indexer/pull/257) (docs):**** DO has NO Singapore/India GPU (NA-only droplets) ? GPU+bucket move to ****GCP asia-south1 (Mumbai)****. ``deploy/gcp/{README,provision.sh,teardown.sh}`` = full runbook + live-resource inventory + switch-off. ****PROVISIONED LIVE**** (project ``gen-lang-client-0306026826``, billing ****Khelsutra Billing**** ``01420D-419EE0-53D83C``): all APIs on · bucket ****gs://khelsutra-rally-corpus`** (asia-south1) · SA **\*\*`rally-worker`**** + key at

```

**~/gcp/rally-worker-sa.json** (local only, gitignored) . **$100` budget `rally-100usd-
guard** (50/90/100%).
49 - ? **FIRST REAL GPU RUN DONE on GCP (2026-06-20) ? report
[#261](https://github.com/avidullu/badminton-highlight-indexer/pull/261)
`docs/DECOUPLED_COMPUTE/DEPLOY_REPORT_GCP_2026-06-20.md`.** GPU quota GRANTED
(`GPUS_ALL_REGIONS`?1, via Cloud Quotas REST API ? **Spot does NOT bypass it, tested**). Spun an
**L4** (`g2-standard-4`, asia-south1-c; T4 was stocked-out in -a; image `common-
cu129-ubuntu-2204-nvidia-580`), ported via `bootstrap --gpu` + the FIXED `gpu_setup.sh`. **M1
? ** `worker infer detector_impl=native` on the smoke clip ? **22 rallies** (gcs in?out);
**highlight reel** stitched ? `gs://khehsutra-rally-
corpus/outputs/05e0e577?/highlight_smoke_l4.mp4`. **M3.5 ? ** native runner on a real cloud GPU;
**cross-GPU parity = sub-pixel (684/1195 frames differ, cuDNN L4-vs-2070) ? gate at the RALLY
level, not CSV bytes** (same-GPU is byte-exact). **M2 ? ** real-extract train validated IN-FLIGHT
(16k+/21k windows extracted on the L4) ? full weights = next session. **VM DELETED** (clean
slate, $0; bucket persists). **Ops Agent fixed** (root cause: project's **logging+monitoring
APIs were OFF** ? now enabled + SA monitoring/logging roles + agent installed).
50 - **BUCKET state** `gs://khehsutra-rally-corpus`: `models/wasb_badminton_best.pth.tar` .
`labels/*.rallies.csv` (7) . `videos/` (smoke + mahadevpura-1 + GX030094_proxy) .
`outputs/05e0e577?/{intervals.csv,report.json,highlight_smoke_l4.mp4}` .
`outputs/parity/native_smoke_l4.csv`.

```

9. assistant (2026-06-21T13:29:58.339Z)

Strong grounding from the handoff. Now let me scout the actual repo surface (packaging-relevant facts) and the key design docs before launching the research workflow. Let me gather structural facts in parallel.

10. user (2026-06-21T13:29:59.354Z)

```

1      # PEP 621 packaging for the badminton/sports highlight indexer (DEFERRED #7).
2      #
3      # Build backend: hatchling. `backend/` is an intentional NAMESPACE package
4      # (no `backend/__init__.py`; mypy.ini sets explicit_package_bases for the same
5      # reason), so the wheel target packages are listed EXPLICITLY below ? hatchling's
6      # auto-detection would otherwise miss the un-__init__'d `backend` root.
7      #
8      # torch/torchvision are deliberately NOT listed in [project.dependencies]:
9      # their CPU/CUDA wheels need an out-of-band index (--index-url / --extra-index-url),
10     # which PEP 621 cannot express. Install them separately ? see [project.optional-
dependencies]
11     # `gpu` (documentation only) and the README. CI installs CPU torch on its own line.
12     [build-system]
13     requires = ["hatchling"]
14     build-backend = "hatchling.build"
15
16     [project]
17     name = "badminton-highlight-indexer"
18     version = "0.1.0"
19     description = "AI-assisted rally segmentation and highlight indexing for racket sports."
20     readme = "README.md"
21     requires-python = ">=3.10"
22     license = { text = "MIT" }
23     authors = [{ name = "Avi Dullu", email = "avi.dullu@gmail.com" }]
24     keywords = ["badminton", "sports", "video", "highlights", "rally", "segmentation"]
25
26     # Runtime deps mirror requirements.txt, EXCEPT torch/torchvision (see header).
27     dependencies = [
28         "fastapi>=0.111.0",
29         "uvicorn>=0.30.1",
30         "opencv-python>=4.9.0.80",
31         "numpy>=1.26.4",
32         "pydantic>=2.7.1",
33         "jinja2>=3.1.4",
34         "python-dotenv>=1.0.0",
35         "google-genai>=2.4.0",

```



```

36     "supervision">=0.21.0,<0.30.0",
37 ]
38
39 [project.optional-dependencies]
40 # Test + lint/type tooling ? matches the CI lanes (.github/workflows/ci.yml).
41 dev = [
42     "pytest",
43     "pytest-cov",
44     "httpx",          # fastapi TestClient transport
45     "ruff",
46     "mypy",
47     "PyJWT[crypto]>=2.8", # M9 edge-auth: Cf-Access JWT verify (RS256) ? tests + CI
48 ]
49
50 # M9 edge-auth (Cloudflare Access). DEFAULT-OFF: only needed when
security.edge_auth_enabled=True.
51 # `jwt` is lazy-imported, so the base runtime works without this; install it for
hosted/multi-tenant.
52 auth = ["PyJWT[crypto]>=2.8"]
53
54 # Per-video observability reports (SOFT dependency: the pipeline degrades to a
55 # no-op if absent). Pinned to its git tag per BACKLOG. The private repo is
56 # fetched over SSH; if you lack access, simply omit this group ? reporting
57 # tests skip via pytest.importorskip. For local dev against a sibling checkout:
58 # pip install -e ../sports-obsreport --no-build-isolation
59 reporting = [
60     "obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3",
61 ]
62
63 # `gpu` is DOCUMENTATION ONLY ? it is intentionally empty. torch/torchvision
64 # CUDA wheels require `--index-url`/`--extra-index-url`, which PEP 621 cannot
65 # express, so they are NOT declarable here. Install them out-of-band instead:
66 # # CUDA 12.4 box:
67 # pip install torch>=2.12.0 torchvision>=0.27.0 --index-url
https://download.pytorch.org/whl/cu124
68 # # CPU-only box (also what CI uses):
69 # pip install torch torchvision --index-url https://download.pytorch.org/whl/cpu
70 gpu = []
71
72 # Storage backends (docs/DECOUPLED_COMPUTE/). The default `local` backend needs NONE of
73 # these. Cloud SDKs are imported lazily, so `import backend.storage` works without them;
74 # install the extra for the bucket you use. (boto3's S3 client covers AWS + Cloudflare
R2 +
75 # DO Spaces + MinIO via endpoint_url.) `gdrive` needs NO extra ? it reads a locally
76 # MOUNTED Google Drive (Google Drive for Desktop) as plain files; see
backend/storage/gdrive.py.
77 storage-s3 = ["boto3>=1.34"]
78 storage-gcs = ["google-cloud-storage>=2.16"]
79
80 [project.scripts]
81 # `indexer --server`, `indexer --help`, etc. -> backend/main.py::main()
82 indexer = "backend.main:main"
83
84 [tool.hatch.metadata]
85 # The `reporting` extra pins obsreport to a git tag (a PEP 508 *direct
86 # reference*); hatchling requires opting in to allow those in metadata.
87 allow-direct-references = true
88
89 [tool.hatch.build.targets.wheel]
90 # `backend` is a namespace package (no __init__.py); list both import roots
91 # explicitly so an editable/wheel install exposes `import backend` and
92 # `import training` regardless of cwd.
93 packages = ["backend", "training"]
94

```

11. user (2026-06-21T13:30:02.515Z)

=== TOP-LEVEL ===

CLAUDE.md
README.md
__pycache__
artifacts
backend
backend.db
config.json
deploy
docs
eval_baselines
indexer.db
mypy.ini
output
pyproject.toml
pytest.ini
rally_index.db
requirements.txt
ruff.toml
run_all_tests.py
run_eval.py
run_regression.py
scratch
scripts
tests
tools
training
videos

=== backend/ ===

__pycache__
annotations
api
assets
audio
billing.py
compute
config
detectors
eval
features
infrastructure
main.py
pipeline
providers
reporting
results.py
sports
stitcher
storage
tools
utils
worker

=== deploy/ ===

cloudrun
digitalocean
gcp

=== deploy/cloudrun ===

Dockerfile
README.md

=== deploy/gcp ===

README.md
provision.sh
teardown.sh

=== docs/ (top) ===

ABLATION_LAYER.md
ANALYZERS_AND_RECONCILERS.md
BACKLOG.md
CODE_AUDIT_AND_TEST_HARDENING.md
CODE_MAP.md
COMMERCIALIZATION.md
COMPUTE_DECOUPLED_SERVING
DEFERRED_HARDENING_PLAN.md
DESKTOP_APP_PLAN.md
DOC_STATUS.md
EVALUATION.md
EXECUTION_POLICY.md
EXPERIMENT_HARNESS.md
GOLDEN_REGRESSION_FIXTURES.md
HARDENING_LOOP_HANDOFF.md
HOSTING_PLAN.md
HOW_RALLY_DETECTION_WORKS.md
I18N_PLAN.md
KHELSUTRA_PRODUCT_OVERVIEW.md
MULTI_SIGNAL_FUSION_PLAN.md
NEXT_STEPS.md
OWNED_MODEL_IMPLEMENTATION_PLAN.md
PERSONALIZATION_PLAN.md
PER_VIDEO_WORKSPACE.md
PLATFORM_ARCHITECTURE.md
POST_166_WORK_ITEMS_MILESTONES_AND_RISK_ASSESSMENT.md
PROJECT_DOC_TEMPLATE.md
PROMINENT_COURT_DETECTION.md
QUALITY_ITERATIONS.md
R1_MILESTONE_LAUNCH_REPORT.md
R2_EVAL_METRIC_DESIGN.md
RALLY_DETECTION_GAP_CLOSURE.md
RALLY_DETECTION_QUALITY_REPORT.md
RALLY_QUALITY_RESEARCH.md
README.md
REAL_FOOTAGE_VALIDATION_REPORT.md
ROADMAP.md
SETUP_NEW_MACHINE.md
STORAGE_SHARING_MODEL.md
TRACKNET_WSL_SETUP.md
UI_ALPHA_EXECUTION_PLAN.md
UI_PROPOSAL.md
VIDEO_LOCALITY_MODEL.md
archives
data_pipeline
mockups

12. user (2026-06-21T13:30:05.642Z)

=== console_scripts / entry points search ===

pyproject.toml:80:[project.scripts]

=== how server/worker are invoked (docs) ===

README.md:93:python -m backend.main --server --port 8000
README.md:101:python -m backend.main --video-path "C:/path/to/my_video.mp4"
README.md:107:python -m backend.main --video-path "C:/path/to/my_video.mp4" --segmenter
yolo_hybrid
README.md:108:python -m backend.main --video-path "C:/path/to/my_video.mp4" --segmenter gemini
README.md:109:python -m backend.main --video-path "C:/path/to/my_video.mp4" --segmenter motion

```

README.md:115:python -m backend.main --video-path "C:/path/to/my_video.mp4" --debug-clip 125.5
README.md:122:python -m backend.main --video-path "C:/path/to/my_video.mp4" \
README.md:249:python -m backend.main --video-path "C:\Users\avidu\Videos\LargeTestVideo.MP4"
--max-frames 300
README.md:263: python -m backend.main --video-path "C:/path/to/video.mp4" --segmenter
yolo_hybrid
README.md:266: python -m backend.main --video-path "C:/path/to/video.mp4" --segmenter gemini
README.md:269: python -m backend.main --video-path "C:/path/to/video.mp4" --segmenter motion
docs/CODE_MAP.md:37:| **Committed regression fixture** (video-free golden) |
`eval_baselines/fixtures/<slug>/` | Tracked, **LF-pinned** (`.gitattributes`). Synthetic Tier-0
(or owner-cleared, de-attributed real). Drives the `golden_fixtures` characterization gate with
**no video/GPU/DB**. See `docs/GOLDEN_REGRESSION_FIXTURES.md`; mint via `python -m
backend.eval.golden_fixtures --freeze-synthetic` / `--freeze-real` (refusal-gated). |

```

```

=== ffmpeg usage sites ===
backend/audio/__pycache__/hit_detector.cpython-313.pyc
backend/audio/__pycache__/__init__.cpython-313.pyc
backend/detectors/__pycache__/wasb_infer.cpython-313.pyc
backend/eval/fusion_audio.py
backend/eval/__pycache__/fusion_audio.cpython-313.pyc
backend/main.py
backend/pipeline/audio/hit_detector.py
backend/pipeline/audio/__init__.py
backend/pipeline/audio/__pycache__/hit_detector.cpython-313.pyc
backend/pipeline/audio/__pycache__/__init__.cpython-313.pyc
backend/pipeline/detectors/wasb_infer.py
backend/pipeline/detectors/__pycache__/wasb_infer.cpython-313.pyc
backend/pipeline/segmenters/gemini.py
backend/pipeline/segmenters/yolo_hybrid.py
backend/pipeline/segmenters/__pycache__/gemini.cpython-313.pyc
backend/pipeline/stitcher/compiler.py
backend/pipeline/stitcher/__pycache__/compiler.cpython-313.pyc
backend/pipeline/validators/base.py
backend/pipeline/validators/format.py
backend/pipeline/validators/pts_continuity.py

```

13. assistant (2026-06-21T13:32:34.274Z)

I now have the packaging surface well-grounded: `pyproject.toml` already ships a `indexer = backend.main:main` console script (hatchling, namespace package), torch/torchvision are deliberately *not* declarable (need `--index-url`), there's a superseded `DESKTOP\_APP\_PLAN.md`, a live `COMPUTE\_DECOUPLED\_SERVING/` track, and the SPA lives in the *sibling* `kheisutra` repo. That's enough to ground an exhaustive research workflow.

Since Ultracode is on and this is a substantial research+design task, I'll launch a multi-phase Workflow: understand the engine's packaging-relevant surface in parallel, research distribution strategies (web-backed, adversarially verified), enumerate the completed CUIs to preserve, then synthesize. I'll author the final plan myself afterward, grounded in the project's conventions.

14. user (2026-06-21T13:32:35.335Z)

Chapter marked: "Batteries-included app plan". Continue your current work.

15. assistant (2026-06-21T13:32:38.290Z)

Now I'll launch the exhaustive research workflow.

16. user (2026-06-21T13:34:48.416Z)

Workflow launched in background. Task ID: w7bxmlxf3
Summary: Research + design a downloadable batteries-included Python app for the badminton highlight indexer (local + cloud inference/training/annotation)
Transcript dir: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\subagents\workflows\wf_a2b97cff-dcd

Script file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\batteries-included-app-plan-wf_a2b97cff-dcd.js

(Edit this file with Write/Edit and re-invoke Workflow with {scriptPath:

"C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\batteries-included-app-plan-wf_a2b97cff-dcd.js"} to iterate without resending the script.)

Run ID: wf_a2b97cff-dcd

To resume after editing the script: Workflow({scriptPath: "C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\batteries-included-app-plan-wf_a2b97cff-dcd.js", resumeFromRunId: "wf_a2b97cff-dcd"}) ? completed agents return cached results.

You will be notified when it completes. Use /workflows to watch live progress.

17. assistant (2026-06-21T13:35:09.777Z)

I've synced (branch is even with `origin/master`; siblings fetched) and grounded myself in the current state, then launched an exhaustive research workflow rather than answering from a partial read. Here's where things stand:

****Done so far****

- ****Repo-freshness**** primary repo even with `origin/master` (`43eccc1`); siblings fetched ? `khelsutra` master at `4525c5d`, `rally-annotator` `d5d59f2`, etc. There's uncommitted ****M9-auth WIP**** in the tree (`backend/api/auth.py`, `models.py`, a new `KHELSUTRA\_PRODUCT\_OVERVIEW.md`) ? likely a sibling session's/owner's work, so I'm treating it strictly read-only and won't touch it.

- ****Grounded**** in the packaging surface: `pyproject.toml` already ships an `indexer = backend.main:main` console script; torch is deliberately out-of-band (`--index-url`); the decoupled-compute profiles (truly-local / cloud-serving / cpu-mock) are merged; the SPA lives in the ***sibling* `khelsutra`** repo; the hard problem is the ****dual torch env**** (main 2.6 vs WASB native 1.11), plus ffmpeg + GCS-hosted weights.

****Running now (background workflow `wf\_a2b97cff`, watch with `/workflows`)**** ? 5 phases, ~28 agents:

1. ****Understand**** (9 parallel readers) ? entrypoints/packaging surface, config+profiles, storage/compute backends, the dual-torch/weights/ffmpeg crux, web-UI bundling, annotation+training, prior appliance/desktop design, ****the completed-CUJ inventory to preserve****, and tests/CI/licensing.
2. ****Research**** (9 web-backed lenses) ? distribution channels (pip/pipx/uv), Docker/Compose, frozen binaries (PyInstaller/Nuitka), the dual-env fix (micromamba/conda-pack/subprocess), weights+ffmpeg bundling, GPU/CPU/MPS cross-platform, launcher UX, ****prior art**** (Ollama/ComfyUI/Frigate/LM Studio), and cloud one-click.
3. ****Verify**** ? each lens's load-bearing claims adversarially refuted (licensing facts, feasibility).
4. ****Critique**** ? completeness critic for gaps.
5. ****Synthesize**** ? architect a tiered packaging plan, options matrix, phased milestones, CUJ-preservation map.

When it lands, I'll write up the detailed plan ? grounded in the project conventions (build on the existing decoupled seam, MIT/Apache/BSD-only, preserve every CUJ, small default-OFF increments) and flagging Claude-doable vs owner-gated work. I'll wait for it to complete rather than poll.

18. user (2026-06-21T14:12:00.096Z)

<task-notification>

<task-id>w7bxmlxf3</task-id>

<tool-use-id>toolu_019fApTT4n1vnyKUNb6Hab5L</tool-use-id>

<output-file>C:\Users\avidu\AppData\Local\Temp\claude\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\tasks\w7bxmlxf3.output</output-file>

<status>completed</status>

<summary>Dynamic workflow "Research + design a downloadable batteries-included Python app for the badminton highlight indexer (local + cloud inference/training/annotation)"

completed</summary>

<result>{"understanding":[{"area":"App launch + current packaging surface (how a user goes from nothing to a running server+worker; what `pip install .` provides; entrypoints; config/DB/output resolution; server-vs-worker process model)","summary":"Today the engine is NOT a downloadable batteries-included app ? it is a source checkout you `pip install -e .` from. The PEP 621 build (hatchling, `pyproject.toml:12-93`) installs the `backend` + `training` namespace packages and registers ONE console script, `indexer = backend.main:main` (`pyproject.toml:80-82`), but deliberately omits torch/torchvision (their CUDA/CPU wheels need `--index-url`, which PEP 621 can't express), so `pip install .` alone yields a torch-less install that cannot run the GPU/yolo_hybrid/trajectory paths. ffmpeg/ffprobe are unmanaged system deps. The intended cold-start path is: (1) install ffmpeg, (2) `pip install -e .` (or `.[dev]`), (3) install torch out-of-band, (4) run `scripts/doctor.py` / `indexer --setup` to GPU-probe + write a default `config.json` + `pip install -r requirements.txt` --user, (5) set `GEMINI\_API\_KEY` in a `.env`, (6) launch. The launcher resolves config.json, .env, the SQLite DB, and output/ all CWD-RELATIVE ? there is no app-data dir, no XDG/AppData home, no env override for their location ? so the running directory IS the implicit state root. Critically, server and worker are NOT separate processes for the normal serving path: the FastAPI server runs an IN-PROCESS ThreadPoolExecutor job worker (`backend.api.job\_worker`), so `python -m backend.main --server` is fully self-contained. `python -m backend.worker {infer|train}` is a SEPARATE, storage-aware (pull?run?push) orchestrator built for the decoupled-compute / Cloud Run path, not a daemon the server talks to. `main()` is overloaded: `--server`, single-shot `--video-path`, `--setup`, `--debug-clip`, `--trim-\*` all dispatch from the one entrypoint, defaulting to server mode when no video path is given (`backend/main.py:557-736`)."}, {"key_facts":[{"fact":"The ONLY console script is `indexer = backend.main:main`; there is no separate `indexer-worker` / `indexer-server` entrypoint. The worker is reached only via `python -m backend.worker`."}, {"evidence":"pyproject.toml:80-82; backend/worker/__main__.py:516-517"}], {"fact":"`pip install .` installs fastapi, uvicorn, opencv-python, numpy, pydantic, jinja2, python-dotenv, google-genai (HARD dep), supervision ? but NOT torch/torchvision (must be installed out-of-band with --index-url). So a bare wheel install cannot run yolo_hybrid/wasb_hybrid/trajectory/GPU paths."}, {"evidence":"pyproject.toml:8-37,63-70; requirements.txt:1-11 (requirements.txt DOES list torch but is only used by doctor.py, not by pip install .)"}], {"fact":"Server and worker are the SAME process for serving: /api/process creates a 'queued' DB job and submits an in-process ThreadPoolExecutor drain (`\_get\_executor().submit(\_drain\_due\_jobs)`); the FastAPI app re-exports the job_worker drain/wake loop. No external worker daemon is required to serve."}, {"evidence":"backend/main.py:80-97,275-315,729-736"}], {"fact":"`python -m backend.worker {infer|train}` is a standalone storage-aware orchestrator (pull URI -> validate/segment -> push outputs/<video_id>/...), composing the same registries as the CLI. It is the decoupled-compute / Cloud Run path, dispatched remotely when compute_target is cloud (M6), NOT a process the --server talks to."}, {"evidence":"backend/worker/__main__.py:1-15,156-259; cmd_infer line 166-168 dispatches remote when get_compute_backend!=None"}], {"fact":"config.json is resolved CWD-relative: DEFAULT_CONFIG_PATH = Path('config.json'); a missing file silently degrades to all-defaults. No env var / app-home override for its location."}, {"evidence":"backend/config/loader.py:30,102-122"}], {"fact":"The SQLite DB path is output_dir-relative: db_path = os.path.join(global_config.output.output_dir, db_name); output_dir defaults to the literal 'output' (CWD-relative), overridable only by the CLI --output-dir flag."}, {"evidence":"backend/main.py:636-640; backend/config/models.py:230-235 (OutputConfig output_dir='output', db_name='sports_indexer.db')"}], {"fact":".env is loaded CWD-relative at import (`load\_dotenv()` with no path). GEMINI_API_KEY is read straight from os.environ by the gemini/yolo_hybrid/trajectory segmenters; absent key => silent 0-segment run flagged in the report."}, {"evidence":"backend/main.py:8-10; backend/pipeline/segmenters/gemini.py:90, yolo_hybrid.py:80, trajectory_hybrid.py:169; README.md:103"}], {"fact":"The server mounts and serves a zero-build static UI from the CWD-relative path 'backend/api/static' (os.makedirs at import), serving index.html/style.css/app.js. This is the legacy in-engine dashboard ? NOT the khelusutra Vite/React SPA ? and the mount path is hardcoded relative to CWD."}, {"evidence":"backend/main.py:64-74; backend/api/static/{index.html,style.css,app.js} exist"}], {"fact":"`indexer --setup` shells out to scripts/doctor.py (path resolved from __file__, cwd-independent), which GPU-probes via nvidia-smi, writes a default config.json via save_default_config(), and runs `pip install --user -r requirements.txt` with the matching torch --extra-index-url (cu124 if GPU else cpu; macOS uses default wheels)."}, {"evidence":"backend/main.py:575-587; scripts/doctor.py:15-42,44-66,129-133"}], {"fact":"requires-python >=3.10; build backend hatchling; backend is an intentional namespace package (no backend/__init__.py) so wheel packages are listed explicitly as ['backend','training']."}, {"evidence":"pyproject.toml:3-6,21,89-93"}], {"fact":"Crash recovery on server/CLI startup sweeps jobs left 'queued'/'running' by a prior process to 'failed' (the

executor is in-process, so they're dead) ? confirms the single-process serving model.", "evidence": "backend/main.py:642-646; comment 'the executor is in-process'", {"fact": "run_all_tests.py is a thin wrapper over `python -m pytest` (defaults to -q), fully mocked/offline, with obsreport reporting tests skipping if the optional dep is absent.", "evidence": "run_all_tests.py:36-55"}, {"fact": "config save (save_default_config) and config.json discovery both anchor on Path('config.json') relative to CWD; doctor.py instead writes to REPO_ROOT/config.json (os.path.dirname x2 of __file__), so doctor and a launched server can disagree on which config.json is authoritative if launched from a different CWD.", "evidence": "backend/config/loader.py:196-206; scripts/doctor.py:9,57"}], "packaging_implications": ["A batteries-included app cannot ship as a plain wheel: torch/torchvision are out-of-band by design (pyproject.toml:8-11,63-70). The packaging layer must provide a post-install/first-run torch bootstrap (the logic already exists in scripts/doctor.py:15-42 ? GPU-probe -> pick cu124/cpu index) or ship a platform-pinned bundle. This is THE first hard problem and is independent of the WASB conda-env split.", "ffmpeg/ffprobe are unmanaged system deps checked only by doctor.py (scripts/doctor.py:84-109). A downloadable app must bundle static ffmpeg/ffprobe binaries (license-clean: ffmpeg LGPL/GPL ? verify which build) or vendor a pip-installable static-ffmpeg, and point the engine at them, because today everything assumes they are on PATH.", "ALL mutable state (config.json, .env, SQLite DB, output/, the static UI mount) is CWD-relative with NO app-home/XDG/AppData concept (backend/config/loader.py:30; backend/main.py:64-66,636-640). A packaged app run from an arbitrary launch dir (e.g. an OS app menu) would scatter or fail to find state. Packaging must introduce a resolved app-data root (config + DB + output + cache) and thread it through OutputConfig.output_dir, DEFAULT_CONFIG_PATH, load_dotenv path, and the static mount.", "The single console entrypoint `indexer` already gives a clean app launch surface (pyproject.toml:80-82). A packaged launcher can default to `indexer --server` and rely on the in-process job worker (backend/main.py:729-736) ? NO separate worker pro ... (truncated 455580 chars, full result in C:\\Users\\avidu\\AppData\\Local\\Temp\\claude\\C--Users-avidu-Projects-badminton-highlight-indexer\\c7679509-ee2b-4443-b240-0b3b1b1a7291\\tasks\\w7bxmlxf3.output)</result><usage><agent_count>29</agent_count><subagent_tokens>3061591</subagent_tokens><tool_uses>806</tool_uses><duration_ms>2231633</duration_ms></usage></task-notification>

19. assistant (2026-06-21T14:12:06.863Z)

The workflow completed (29 agents, ~3M tokens, 806 tool uses). The result is large ? let me read the full output, focusing on the synthesis and critic sections first.

20. user (2026-06-21T14:12:09.211Z)

```
=== SIZE ===
506396 C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c7679509-ee2b-4443-b240-0b3b1b1a7291/tasks/w7bxmlxf3.output
=== byte offset of synthesis ===
474347:"synthesis"
=== byte offset of critic ===
453371:"critic"
```

21. user (2026-06-21T14:12:16.759Z)

```
TOP KEYS: ['summary', 'agentCount', 'logs', 'result']
understanding areas: 0
research topics: 0
```

```
===== CRITIC =====
{}
```

22. user (2026-06-21T14:12:24.443Z)

```
RESULT KEYS: ['understanding', 'research', 'critic', 'synthesis']
  understanding: list 9
  research: list 9
  critic: dict 5
  synthesis: dict 9
```

saved parsed result

23. user (2026-06-21T14:12:29.737Z)

===== SYNTHESIS =====

--- EXECUTIVE SUMMARY ---

Ship the engine as ONE pip-installable artifact whose behavior is selected entirely by the already-merged decoupled-compute seam (deployment.profile / compute_target / StorageBackend / ComputeBackend), fronted by a thin friendly launcher and a Windows installer ? NOT a frozen monolith and NOT an Electron/Tauri shell. The single hardest constraint forces the scope: torch is deliberately out of pyproject (needs --index-url), and native WASB requires a SECOND, incompatible torch env (conda py3.8 + torch 1.11+cu113) shelled out via \$WASB_PYTHON, so no single venv/freeze can deliver the GPU-quality path. The honest, shippable v1 is therefore a LIGHT tier (Gemini cloud-brain + motion + cpu-mock stub) that needs zero torch, zero GPU, and preserves CUJ-1.5 and CUJ-8; the native-WASB GPU path (CUJ-6/7) ships ONLY to the owner / a separately-provisioned env until two blockers clear: WASB-C3 weight provenance (gate_verdict in training/gen0/harness.py:177 hardcodes ship=False) and a reproducible truly-local native config. Five engine prerequisites must land BEFORE any packaging push, each as a small default-OFF PR: (1) a resolved APP-DATA HOME threaded through config/.env/DB/output/static (today everything is CWD-relative ? loader.py:30 reads Path('config.json') from CWD while scripts/doctor.py:57 writes it to REPO_ROOT, a silent data-loss footgun); (2) redact GET /api/config which today returns the whole AppConfig object unauthenticated (backend/main.py:144-146); (3) make the bind address and the shipped default profile config-driven (uvicorn hardcodes host='127.0.0.1' at main.py:756); (4) declare scipy (serving-required via backend/features/rally_seq.py, only transitive today) and add a license-scan + built-artifact CI lane; (5) build the three net-new endpoints the UI/wizard depend on (/api/capabilities, /api/videos list, /api/reels) ? confirmed absent. Distribution leads with a generated install script over uv (auto-installs Python + the right torch index) plus an Inno Setup Windows installer; a Docker image (the existing deploy/cloudrun/Dockerfile, micromamba-ified) serves the cloud/server profile and the BYO-GPU path. All cloud, edge-auth, cost-ledger, and training features stay default-OFF so the local single-user app is byte-identical to the verified CUJs.

--- RECOMMENDED ARCHITECTURE ---

ONE build, profile-selected. Ship a single pip wheel (the existing hatchling build: console script `indexer = backend.main:main`, packages=[backend, training]) plus the cross-repo SPA built into it as package data. The pure DETECT->WINDOW->STITCH core never branches on profile (design D1); a single `deployment.profile` switch (backend/config/presets.py PROFILE_PRESETS: truly-local / cloud-serving / cpu-mock) applies a coherent preset bundle, and `get\_compute\_backend(config)` (backend/compute/__init__.py) returns CloudRunCompute only for compute_target=='cloud_run' else None=in-process. This is the seam to BUILD ON, not reinvent ? local vs cloud is genuinely one config flip, and the dispatched Cloud Run Job runs the SAME image forced inline. The app defaults to the truly-local preset (skip_ai_handoff handled per tier) so the out-of-box experience is zero-cloud unless the user opts in.

HOW ONE BUILD RUNS LOCAL + CLOUD: the local FastAPI server (with its in-process ThreadPoolExecutor(max_workers=1) job worker, backend/api/job_worker.py:37) is the control plane in every profile. truly-local runs detect/stitch in-process. cloud-serving flips compute_target=cloud_run so the local control plane dispatches a Cloud Run GPU Job and bucket-polls a done-marker; storage.backend=gcs routes I/O. cpu-mock gives a \$0 deterministic self-test on first run. The user toggles profile in Settings (rendered from a NEW /api/capabilities endpoint) with explicit confirm (D15 ? cloud spend never silent) and an INR cost estimate from backend/billing.py before any cloud job.

HOW THE UI IS BUNDLED: build khelsutra/web (`npm ci && npm run build`) at package time and vendor web/dist/ into the wheel via a hatchling force-include mapping to a package dir (e.g. backend/ui/), served by FastAPI with an SPA history-fallback route. The SPA already calls same-origin `const BASE='/api/'` (client.ts:7) so ZERO SPA code change is needed for single-origin serving ? this collapses the appliance's reverse-proxy topology into one FastAPI process. Retire the legacy CWD-relative backend/api/static mount (main.py:71-72) and replace the hardcoded relative paths with importlib.resources package-relative resolution. dist/ is gitignored, so the build must produce it and record the bundle hash + an /api/version handshake + a contract test pinning SPA->engine.

HOW THE DUAL-TORCH / WEIGHTS / FFMPEG PROBLEMS ARE SOLVED: (a) DUAL-TORCH ? do NOT try to host

both ABIs in one venv or freeze. v1 LIGHT tier carries no torch at all (Gemini/motion/stub). The GPU/WASB HEAVY tier is a SEPARATELY provisioned second environment: short-term via a micromamba+conda-forge env (license-clean, sidesteps the Anaconda-defaults ToS the current deploy/cloudrun/Dockerfile:35 still trips) reached by \$WASB_PYTHON exactly as native_wasb_runner.py:89 already does; the strategically correct long-term fix is to EXPORT WASB to ONNX (the DetectorRunner ABC already names an ONNXRunner writing the identical Frame,Visibility,X,Y CSV) so the WASB path needs neither torch 1.11 nor 2.x ? onnxruntime (MIT, ~tens of MB, CPU+CUDA+DirectML+CoreML) collapses the dual-env entirely and is owner-gated on a byte-parity re-validation. (b) WEIGHTS ? never bake wasb_badminton_best.pth.tar into the artifact (size + WASB-C3 provenance UNRESOLVED). Fetch on first run with a pinned SHA-256, resumable, into the app-data cache, setting \$WASB_WEIGHTS. CRITICAL: until counsel clears C3, the public download must NOT auto-fetch these weights ? the WEIGHTS path is owner-only in v1. (c) FFMPEG ? do NOT redistribute ffmpeg (the current encoder is GPL libx264 at compiler.py:253 and video.py:42). Route every ffmpeg/ffprobe call through ONE resolver helper (config path / env / shutil.which / fail-loud), and have first-run setup (extending scripts/doctor.py, which already prints the winget command) detect-and-offer `winget install Gyan.FFmpeg`. Keep an LGPL-ffmpeg-bundle (BtbN,openh264/NVENC) as a documented fallback that also requires swapping the libx264 default ? tracked, not v1.

TIERS: LIGHT (v1, default) ? pip/uv install or Windows installer, no torch, Gemini/motion/cpu-mock, full reel CUJ on any CPU box. HEAVY/local-GPU (owner + power user) ? adds the micromamba WASB env + weight fetch, gated on C3 + reproducible config + an ONNX milestone. CLOUD/server ? the Docker image for a GCP/Cloud Run or BYO-GPU-VM deploy behind Cloudflare Access edge-auth.

--- DISTRIBUTION TIERS ---

? LIGHT (v1 default, recommended) | audience: Non-tech India coach/parent on Windows; any CPU-only machine; the broad beachhead

contents: pip wheel (indexer console script) + bundled SPA + Gemini(google-genai, hard dep)/motion/cpu-mock segmenters. NO torch, NO GPU, NO WASB weights. ffmpeg resolved via winget-offer at first run. Defaults to truly-local profile; Gemini needs a user-supplied key (consent copy that <=1280px chunks leave the machine, never a training source).

install_ux: Inno Setup .exe (Start-Menu shortcut -> friendly launcher -> auto-opens 127.0.0.1:8000) as the primary non-tech path; a generated install script over uv (auto-downloads Python 3.10+, isolates the venv, installs the app) as the cross-platform/power-user path. First run: cpu-mock self-test proves the install, then a 3-step in-UI wizard (key/compute/videos).

? HEAVY / local-GPU (owner + opt-in power user) | audience: Owner with the golden corpus + a CUDA box; later, GPU-equipped power users once blockers clear

contents: LIGHT + a SEPARATELY provisioned WASB env (micromamba+conda-forge, py3.8/torch1.11+cu113) reached via \$WASB_PYTHON; main-env torch installed out-of-band via uv --torch-backend=auto; WASB weights fetched from GCS (SHA-256-pinned) at first run. GATED on WASB-C3 counsel sign-off + a reproducible native config; the ONNX-export milestone supersedes this env when it lands.

install_ux: An opt-in 'Enable local GPU detection' step in the launcher/installer that provisions the second env with a loud resumable progress UI (multi-GB torch wheels on Indian bandwidth) and a post-install torch.cuda.is_available() validation; falls back to LIGHT with a 'GPU not available, running cloud/CPU' banner.

? CLOUD / server | audience: GCP deploy (Cloud Run GPU Jobs, per ADR-013) and BYO-GPU-VM (deploy/gcp/create_gpu_vm.sh); the hosted-service or self-host-on-a-box operator

contents: The deploy/cloudrun/Dockerfile worker image (Miniconda -> micromamba+conda-forge to clear the Anaconda ToS), bundling ffmpeg + the dual-env stack; runs cloud-serving (compute_target=cloud_run, storage=gcs) behind Cloudflare Access. Closes the per-video M8 cost ledger when worker usage-emission is wired.

install_ux: `gcloud run jobs create` per deploy/cloudrun/README.md, or create_gpu_vm.sh/teardown.sh for a warm box. A secondary 'Deploy to Render' README button is CPU-only control-plane only (Render/Railway have NO GPU; Fly GPUs deprecated) and must dispatch GPU work back to GCP. Server bind must become 0.0.0.0 (config-gated) + edge-auth mandatory.

24. assistant (2026-06-21T14:12:34.177Z)

Excellent, grounded synthesis. Let me pull the remaining synthesis sections plus the critic and research recommendations.

25. user (2026-06-21T14:12:37.615Z)

=== OPTIONS MATRIX ===

? v1 scope: which tier ships first

options: ['All-in-one GPU+WASB bundle', 'LIGHT (Gemini/motion/cpu-mock) first, HEAVY/WASB owner-only', 'Cloud-only hosted service']

REC: LIGHT first; HEAVY owner-only behind C3 + reproducible-config gates

why: The dual-torch reality + WASB-C3 ship=False (harness.py:177) + the unverified reproducible native config make an all-in-one GPU public download unshippable and likely licensing-violating. LIGHT preserves CUJ-1..5/8 with zero torch and is honestly shippable now. (Tension: this temporarily drops the verified WASB CUJ-6/7 from the DEFAULT ? surface it explicitly, do not pretend every CUJ ships in v1.)

? Primary distribution channel

options: ['PyPI wheel + pip', 'pipx', 'uv tool install + generated script', 'Inno Setup installer', 'PyInstaller frozen onefile']

REC: Inno Setup installer (non-tech) backed by a uv-based generated script (power/cross-platform); reject PyInstaller

why: uv is the only channel that auto-installs Python + routes the torch --index-url; an Inno shortcut is the non-scary 'open the app' affordance. PyInstaller+CUDA is 3-5GB, can't split CUDA libs, and cannot host the second WASB env. CAVEAT: uv --torch-backend=auto is uv-pip-only (not uv tool install) ? the script must run uv pip into the venv and validate torch.cuda post-install; verify against a real uv before relying on it.

? Dual-torch WASB delivery

options: ['Both torches in one venv (impossible)', 'conda-pack relocatable tarball', 'micromamba+conda-forge env on first run', 'ONNX export (no torch on WASB path)', 'Docker-only']

REC: micromamba+conda-forge now as the bridge; ONNX export as the strategic fix (owner-gated re-validation)

why: Two incompatible ABIs cannot share a venv or a freeze. micromamba is BSD and sidesteps the Anaconda-defaults ToS the current Dockerfile:35 still trips. ONNX (onnxruntime MIT) dissolves the dual-env, ships tens-of-MB not 5GB, and gives CPU+CUDA+DirectML+CoreML ? but needs a byte-parity gate on a real GPU (owner-side, offline CI can't prove numbers).

? ffmpeg delivery

options: ['System winget install + resolver', 'Bundle imageio-ffmpeg', 'Redistribute LGPL static build']

REC: Winget-install + one resolver helper; LGPL-bundle as documented fallback

why: Not redistributing ffmpeg sidesteps the GPL-vs-LGPL landmine entirely (current encoder is GPL libx264). imageio-ffmpeg lacks ffprobe (which the engine requires) and its binary is effectively GPL. An LGPL bundle would force swapping libx264 -> openh264/NVENC at BOTH sites (compiler.py:253, video.py:42).

? Launch UX wrapper

options: ['Bare .bat + webbrowser.open', 'pystray system-tray launcher', 'Electron/Tauri embedded-webview shell']

REC: Friendly launcher that health-polls then opens the browser; REJECT pystray (LGPL-3.0) for an MIT tray lib or stdlib-only; reject Electron/Tauri

why: D10 locks 'self-hosted webserver, NOT a desktop app'. A launcher that opens the user's own browser at 127.0.0.1:8000 honors that; an embedded-webview shell IS a desktop app and adds WebView2/Chromium friction + presumes a PyInstaller freeze. pystray is LGPL-3.0 ? forbidden by the guardrail; pick infi.systray (or equivalent MIT) or ship the bare launcher. Get owner re-ratification in writing.

? Cloud target

options: ['Cloud Run GPU Jobs', 'Render/Railway', 'Fly.io GPU', 'BYO GPU VM']

REC: Cloud Run GPU Jobs (primary, ADR-013) + BYO-GPU-VM (warm path); Render only as CPU control-plane

why: Render/Railway have no GPU; Fly GPUs are deprecated. Cloud Run L4 is per-second, scale-to-zero, GCP-co-located with the bucket + cost ledger, and the CloudRunCompute dispatch is already built/mock-tested. Render/Railway pull compute off GCP (ADR-013 tension) so only as a CPU front door dispatching back to GCP.

? Edge auth before any non-localhost exposure

options: ['None (localhost only)', 'Cloudflare Access + JWT verify', 'Reverse-proxy basic-auth']

REC: Cloudflare Access (cloudflared tunnel) + in-app PyJWT signature verify + Cloud Run ingress lock

why: The engine has zero app auth by default and the GET /api/config dump leaks the full config; the ONLY safe default is bind-127.0.0.1. For exposure, D9 mandates BOTH edge gate AND in-app JWT signature verify (header-trust alone is spoofable). auth.py is already wired default-OFF at main.py:294; M9 is counsel-gated for public signup.

=== PHASED MILESTONES ===

? P0 ? Engine prerequisites (Claude-doable, default-OFF, before any packaging)
(owner_gated=False)

goal: Make the engine safe to run from an arbitrary launch dir and to bundle, with no behavior change when flags are off

- Introduce a resolved APP-DATA HOME (env override RALLY_APP_HOME + OS default %APPDATA%/Khelsutra | ~/.local/share/khelsutra); thread it through loader.DEFAULT_CONFIG_PATH (loader.py:30), the load_dotenv(path=...) call (main.py:10), OutputConfig.output_dir/db path, and the static mount; unify scripts/doctor.py:57 REPO_ROOT write with the server's read so --setup and the server agree
 - Redact GET /api/config (main.py:144-146 returns the whole AppConfig); add a test asserting no credential-shaped field ever serializes
 - Make uvicorn bind (main.py:756) and CORS origin config/profile-driven: default 127.0.0.1, widen to 0.0.0.0 only under an edge-auth profile; add a Host-header/localhost-origin guard against DNS-rebinding to /api/config
 - Declare scipy in [project.dependencies] (used in backend/features/rally_seq.py, only transitive today); add a pip-licenses/reuse CI lane and a built-artifact smoke lane (build, then indexer --help + boot + /api/health)
 - Add a single-instance guard (PID/lockfile) so a second launch can't open the same SQLite DB and have fail_stale_jobs (main.py:664) sweep the first process's live jobs
- exit: All 5 CI lanes green at coverage >=70%; launching `indexer --server` from any CWD finds the same config/DB; /api/config leaks nothing; new lanes pass

? P1 ? UI bundling + net-new endpoints (owner_gated=False)

goal: Serve the real SPA single-origin from the wheel and build the endpoints the wizard/console need

- Cross-repo build step: npm ci && npm run build in khelsutra/web, vendor web/dist/ into the wheel via hatchling force-include -> backend/ui/, served by FastAPI with an SPA history-fallback; retire the legacy backend/api/static mount; record the bundle hash
 - Build GET /api/capabilities (honest hardware/segmenter/ffmpeg/weights probe), GET /api/videos (list; today only single-video get), GET /api/reels/{video_id} (reload-survival)
 - Add an /api/version handshake + a contract test pinning the bundled SPA against the engine API (dist/ is gitignored so there's no auditable record otherwise)
- exit: FastAPI serves the SPA at / with refresh-survival; capabilities/videos/reels return real data; contract test guards SPA<->engine skew

? P2 ? LIGHT-tier installer + first-run UX (owner_gated=False)

goal: A non-tech Windows user downloads, installs, and produces a reel with no terminal

- Inno Setup installer: lay down a uv-managed venv + the app, Start-Menu/desktop shortcut to a friendly launcher (health-poll -> webbrowser.open; free-port fallback; graceful shutdown); winget-offer ffmpeg via the new resolver helper; ship a config.json defaulting to truly-local; budget a code-signing cert for SmartScreen
 - ffmpeg resolver helper wrapping ALL call sites (compiler.py, video.py); first-run cpu-mock self-test
 - In-UI 3-step wizard (rendered from /api/capabilities): Gemini key (paste/skip, written to app-home .env with consent copy) | Compute (Local CPU / Cloud, only what the box can honestly do) | Videos (local folder / gdrive mount). en/kn/hi localized
 - Uninstall manifest: remove env/binaries, prompt to keep videos/reels, scrub the .env secret
 - DB upgrade contract: back up the SQLite file before running the user_version ladder (database.py at v6); refuse to open a DB whose user_version exceeds the binary's max (downgrade guard); add an 'export my data' affordance
- exit: Clean Windows VM -> installer -> reel downloaded, en+hi, with no Python/git/ffmpeg preinstalled; uninstall is clean; upgrade preserves data

? P3 ? HEAVY/local-GPU + cloud profiles (largely owner-gated) (owner_gated=True)
 goal: Provision the WASB second env + cloud dispatch without breaking LIGHT defaults
 - Replace Miniconda+`conda tos accept` (deploy/cloudrun/Dockerfile:33-37) with micromamba+conda-forge; record an ADR; reuse for the local managed WASB env
 - First-run WASB weight fetcher (SHA-256-pinned, resumable) setting \$WASB_WEIGHTS ? GATED: public build must NOT fetch until C3 clears (owner-only flag)
 - 'Push to cloud' Settings toggle: flip to cloud-serving with confirm + INR cost estimate + a per-video cost cap; wire worker usage-emission into the M8 ledger
 - Land M9 edge-auth path verified (already wired default-OFF at main.py:294): JWKS refresh for the ~6-week rotation, Cloud Run ingress lock
 exit: Owner reproduces the native-WASB reel on a fresh GPU box from the packaged HEAVY profile; cloud toggle dispatches a real Cloud Run Job with a correct INR estimate; LIGHT default unchanged

? P4 ? ONNX export + training/annotation decisions (owner_gated=True)
 goal: Dissolve the dual-env and decide BYO-training/annotation scope
 - Export WASB to ONNX (onnxruntime MIT); implement the ABC's ONNXRunner writing the identical CSV; re-pass the owner byte-parity/golden gate on a real GPU
 - Decide training-in-UI: worker train is an offline CLI and in-process jobs don't survive restart ? design a durable long-job runner + progress stream + cancel, and surface gate_verdict (ship=False) honestly; recommend DEFER from v1
 - Decide annotation: bundle VLC 3.0.x plugin (heavy, version-pinned) vs build the in-browser HTML5 annotator emitting the same rally CSV ? recommend the in-browser annotator, owner-gated
 exit: ONNX runner passes byte-parity; an explicit owner decision recorded for training + annotation v1 scope

26. user (2026-06-21T14:12:46.331Z)

<persisted-output>

Output too large (29.5KB). Full output saved to: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\tool-results\bx2mlt2vm.txt

Preview (first 2KB):

=== CUJ PRESERVATION ===

? CUJ-1 Local Gemini inference -> rallies -> reel (en+hi)
 -> LIGHT tier ships Gemini (google-genai hard dep) + the full
 /api/process->jobs->query->compile->highlights loop on the in-process worker with storage=local; preserved verbatim, only needs the user's Gemini key and a resolved ffmpeg.
 ? CUJ-2 In-UI gs:// bucket ingest
 -> storage.acquire_local_input + the gcs StorageBackend already accept gs:// at
 /api/process; preserved when the user opts into a cloud storage backend (google-cloud-storage as an optional extra + ADC).
 ? CUJ-3 Per-rally selection + deterministic query bar
 -> Pure SPA + /api/query + /api/compile; preserved by bundling the built khelsutra/web SPA single-origin (BASE='/api' unchanged) and the contract test guarding the endpoints it calls.
 ? CUJ-4 Reel download-to-local
 -> /api/compile + GET /api/highlights/{video_id}/{filename} unchanged; the resolved ffmpeg (winget) provides the encoder; app-home output dir keeps reels findable across launches.
 ? CUJ-5 Trilingual SPA (en+kn+hi), offline
 -> Catalogs + OFL Noto fonts are compiled into the bundle (no CDN); vendoring web/dist/ keeps it air-gap renderable. (Gate: kn/hi native-speaker review still pending before non-test users.)
 ? CUJ-6/7 Worker GCS->GPU WASB->GCS + cold-start dual-env
 -> Preserved on the HEAVY/cloud tiers ONLY (owner + provisioned GPU), via the micromamba WASB env / Cloud Run image ? NOT in the LIGHT default. This is the one CUJ deliberately absent from the v1 default; flagged explicitly, gated on WASB-C3 + a reproducible native config.
 ? CUJ-8 CLI single-shot inference (indexer --video-path --segmenter ...)
 -> The console script and module entrypoints are untouched by packaging; available in every tier (segmenter set limited by tier capabilities).
 ? CUJ-9 Annotate -> train (PARTIAL today)

...

</persisted-output>

27. user (2026-06-21T14:12:49.766Z)

```

1      === CUJ PRESERVATION ===
2      ? CUJ-1 Local Gemini inference -> rallies -> reel (en+hi)
3      -> LIGHT tier ships Gemini (google-genai hard dep) + the full
/api/process->jobs->query->compile->highlights loop on the in-process worker with storage=local;
preserved verbatim, only needs the user's Gemini key and a resolved ffmpeg.
4      ? CUJ-2 In-UI gs:// bucket ingest
5      -> storage.acquire_local_input + the gcs StorageBackend already accept gs:// at
/api/process; preserved when the user opts into a cloud storage backend (google-cloud-storage as
an optional extra + ADC).
6      ? CUJ-3 Per-rally selection + deterministic query bar
7      -> Pure SPA + /api/query + /api/compile; preserved by bundling the built
khelsutra/web SPA single-origin (BASE='/api' unchanged) and the contract test guarding the
endpoints it calls.
8      ? CUJ-4 Reel download-to-local
9      -> /api/compile + GET /api/highlights/{video_id}/{filename} unchanged; the resolved
ffmpeg (winget) provides the encoder; app-home output dir keeps reels findable across launches.
10     ? CUJ-5 Trilingual SPA (en+kn+hi), offline
11     -> Catalogs + OFL Noto fonts are compiled into the bundle (no CDN); vendoring
web/dist/ keeps it air-gap renderable. (Gate: kn/hi native-speaker review still pending before
non-test users.)
12     ? CUJ-6/7 Worker GCS->GPU WASB->GCS + cold-start dual-env
13     -> Preserved on the HEAVY/cloud tiers ONLY (owner + provisioned GPU), via the
micromamba WASB env / Cloud Run image ? NOT in the LIGHT default. This is the one CUJ
deliberately absent from the v1 default; flagged explicitly, gated on WASB-C3 + a reproducible
native config.
14     ? CUJ-8 CLI single-shot inference (indexer --video-path --segmenter ...)
15     -> The console script and module entypoints are untouched by packaging; available
in every tier (segmenter set limited by tier capabilities).
16     ? CUJ-9 Annotate -> train (PARTIAL today)
17     -> The worker train CLI + rally-annotator CSV bridge remain runnable; honesty fences
(synth=mock, gate_verdict ship=False) must surface in any UI. Recommended DEFERRED from v1 UI
per P4; never advertised as a closed loop.
18
19     === RISKS & MITIGATIONS ===
20     ? RISK: CWD-relative state (config/.env/DB/output): launching from the OS menu
silently creates an empty config+DB and the user 'loses' their index; doctor.py:57 (REPO_ROOT)
and the server (CWD) disagree on config.json
21     MIT: P0 app-data home threaded through every path + doctor/server unification; this
is the sequencing prerequisite for uninstall, auto-update, and migration
22     ? RISK: Auto-fetching WASB weights into a PUBLIC download while WASB-C3 provenance is
unresolved violates the MIT/Apache/BSD-only weights guardrail (gate_verdict hardcodes
ship=False)
23     MIT: v1 = LIGHT only; weight fetch is an owner-only flag until counsel clears C3;
record C3 resolution as the explicit gate on the HEAVY public path
24     ? RISK: GET /api/config returns the entire AppConfig unauthenticated (main.py:144) ->
any future credential field + a DNS-rebinding hit to 127.0.0.1:8000 leaks the Gemini key
25     MIT: P0 redaction + a serialization test + a Host-header/localhost-origin guard;
keep secrets in app-home .env (or OS keychain), never in config
26     ? RISK: uv --torch-backend=auto is uv-pip-only (not uv tool install) and the cross-
mode 'uv pip into the tool venv' hack is unverified against a real uv; the lead channel's
central mechanic could fail
27     MIT: Verify against a pinned uv version in CI; the installer must run uv pip into
the managed venv and validate torch.cuda.is_available() post-install, falling back to CPU/LIGHT
28     ? RISK: DB migration on auto-update: opening an older DB with newer code (or
downgrading onto a newer user_version) corrupts or fails silently
29     MIT: P2 back-up-before-migrate, a downgrade guard (refuse user_version > binary
max), and an 'export my data' affordance wired into the auto-update flow
30     ? RISK: Reproducible native-WASB config drift (prior cold-start yielded 0 rallies); a
'native' profile could ship a 0-rally out-of-box
31     MIT: Owner reproduces the 22-rally run from the packaged HEAVY profile as a P3 exit
criterion before any non-owner exposure; bake the exact validated config (incl. skip_ai_handoff)
into the preset
32     ? RISK: SPA<->engine version skew: bundled SPA calls endpoints a stale engine lacks
(capabilities/videos/reels are net-new); dist/ gitignored so no auditable record
33     MIT: P1 /api/version handshake + contract test + recorded bundle hash; build SPA at

```

package time pinned to the engine commit

34 ? RISK: pystray is LGPL-3.0 (guardrail forbids); Windows SmartScreen/AV quarantines unsigned launchers; MAX_PATH + non-ASCII usernames break conda/uv paths

35 MIT: Use an MIT tray lib or stdlib webbrowser-only; budget a code-signing cert; enable Windows long-path support + validate the app-home path is ASCII/short, warn early

36 ? RISK: Built-artifact behavior unverified: CI tests source only; hidden imports, namespace-package resolution (no backend/___init__.py), missing data files (fonts) surface only on the user's machine

37 MIT: P0 built-artifact smoke CI lane (build -> indexer --help -> boot -> /api/health); exclude bootstrap code from the coverage floor

38 ? RISK: Training/long-job UX: in-process jobs don't survive restart (fail_stale_jobs) so a multi-hour GPU train can't run as an ephemeral job

39 MIT: Defer training from v1 UI (P4); if added, build a durable out-of-process runner with progress stream + cancel and surface gate_verdict ship=False honestly

40

41 === LICENSING NOTES ===

42 ? WASB-C3 is the dominant blocker: the CODE is MIT but the badminton CHECKPOINT provenance is UNRESOLVED for commercial sharing (Dockerfile:14-16; gate_verdict hardcodes ship=False at training/gen0/harness.py:177). Do NOT conflate MIT-code with cleared-weights ? the public download must not auto-fetch the weights until counsel signs off.

43 ? ffmpeg: the current encoder is GPL libx264 (compiler.py:253, video.py:42). The license-clean path is NOT to redistribute ffmpeg (winget-install + a resolver). Any bundled ffmpeg must be an LGPL build (Btbn,openh264/NVENC) AND requires swapping the libx264 default at both call sites; only ffprobe-bearing full builds satisfy the engine.

44 ? Anaconda ToS: deploy/cloudrun/Dockerfile still uses Miniconda + `conda tos accept` against pkgs/main + pkgs/r (lines 33-37) ? the >=200-employee ToS exposure. Replace with micromamba + conda-forge (BSD, no ToS) for both the cloud image and any local managed WASB env; record an ADR.

45 ? New packaging deps clearance: uv (Apache-2.0/MIT), Inno Setup (open-source/permissive), onnxruntime (MIT), micromamba (BSD-3), Docker Engine/OCI images (Apache-2.0), boto3/google-cloud-storage/google-cloud-run (Apache-2.0), stdlib webbrowser ? all clean. pystray is LGPL-3.0 and is FORBIDDEN by the guardrail; choose an MIT tray lib. Nuitka 4.x is AGPL-3.0+ (NOT Apache-2.0) ? usable only as a build-time tool (compiled output has a runtime exception), not bundled.

46 ? obsreport (git+ssh private repo) must stay a truly-optional extra never bundled into the public download ? a non-repo user can't resolve it; it already degrades to a no-op.

47 ? Gemini stays inference/serving ONLY, never a training data source; the wizard's key step must say so. cloud-serving free-tier data-for-reels training is adults-only/derived-data/raw-auto-delete (D10), counsel-gated under DPDP/GDPR ? out of v1.

48 ? Add an automated license-scan CI lane (pip-licenses/reuse); the guardrail is human-only today, so a license-incompatible packaging dep could slip in unnoticed.

49

50 === OPEN OWNER DECISIONS ===

51 ? RE-RATIFY D10 vs the launcher+installer+(tray): is a Start-Menu shortcut + friendly launcher + Windows installer acceptable as 'a self-hosted webserver the user spins up', or does any OS-menu affordance cross the 'NOT a desktop app' line? Need it in writing before P2.

52 ? Confirm v1 SCOPE = LIGHT tier (Gemini/motion/cpu-mock, no torch, no WASB) and that CUJ-6/7 (native WASB) is deliberately owner-only until WASB-C3 + reproducible-config gates clear. This is the load-bearing scope call.

53 ? WASB-C3: get counsel sign-off on whether the checkpoint may be redistributed / auto-fetched into a public build at all ? determines if a public HEAVY tier is ever possible vs WASB-on-cloud-only.

54 ? Cloud-serving in a downloaded app = (a) a hosted Khelsutra service (no StorageBackend/ComputeBackend client exists for it today) or (b) BYO-GCP (the only thing built)? D13 'your machine or our cloud' currently overstates (a). Decide before marketing the 'push to cloud' toggle.

55 ? Training in v1 UI: include the worker-train CLI + a durable long-job runner, or defer (recommended) given synth=mock, real=GPU+dual-env, and gate_verdict ship=False?

56 ? Annotation in v1: bundle the external VLC 3.0.x plugin vs build the in-browser HTML5 annotator emitting the same rally CSV (recommended) vs defer entirely ? a stated product goal with no implementation path today.

57 ? kn/hi localization: ship the AI-drafted catalogs to real non-test users before native-speaker review, or gate the trilingual UI on review?

58 ? Shipped DEFAULT config.json / DEPLOYMENT_PROFILE: confirm truly-local (skip_ai_handoff per tier) so the box never surprise-bills on Gemini nor returns 0 rallies on

WASB-without-key.

59 ? Code-signing: approve budget for a Windows Authenticode (and later macOS Developer-ID) cert so non-tech users don't hit SmartScreen/AV prompts.

60

61

62 === CRITIC ===

63

64 [gaps]

65 - APP-DATA HOME + DATA MIGRATION is named as an open question but the plan never RESOLVES it, and it is the single most load-bearing gap. ALL mutable state is CWD-relative (load_dotenv() at backend/main.py:10 takes no path; config.json via DEFAULT_CONFIG_PATH=Path('config.json') at loader.py:32; DB at output.output_dir; static mount at main.py:71-72 'backend/api/static'). doctor.py writes config.json to REPO_ROOT (scripts/doctor.py:54) while the server reads it from CWD ? these DISAGREE the moment the launcher's CWD != REPO_ROOT. No plan exists for (a) picking %APPDATA%/~/.local/share root, (b) THREADING it through DEFAULT_CONFIG_PATH + load_dotenv path + OutputConfig.output_dir + the static mount, or (c) MIGRATING an existing CWD-rooted install into it on upgrade. Without this, every distribution channel ships a footgun where launching from the OS menu silently creates an empty config/DB and the user 'loses' their index.

66 - DB SCHEMA MIGRATION ON UPGRADE is unaddressed. database.py has a real PRAGMA user_version ladder (currently at version 6, migrations at lines 183/210/222/242/290/337). A downloadable app that auto-updates will open an OLDER user's DB with NEWER code ? the plan never covers forward-migration safety, backup-before-migrate, or the DOWNGRADE case (user installs v2 app, downgrades to v1 binary, opens a user_version=7 DB). No 'export/backup my data before upgrade' story exists. This is table stakes for any shippable app and is entirely missing.

67 - SECRETS HANDLING for the Gemini key is hand-waved. GET /api/config (main.py:144-146) returns the ENTIRE global_config object unauthenticated with NO field redaction ? if any key/credential field ever lands in AppConfig it leaks to any localhost process/browser tab. The plan repeatedly says 'key goes in env/.env, never config' but never specifies WHERE the wizard writes it (a plaintext .env in the app-data dir? OS keychain?), how it is protected at rest, or how it is scrubbed on uninstall. 'Store locally' (launcher-ux research) is not a secrets design.

68 - UNINSTALL / CLEANUP is never mentioned anywhere. A batteries-included installer that drops a managed Miniconda/uv env, GCS-fetched WASB weights (hundreds of MB), an LGPL ffmpeg binary, a SQLite DB, output reels, and a plaintext Gemini key has NO uninstall story ? what gets removed, what user data is preserved, whether the secret is scrubbed. Inno Setup/winget need an explicit uninstaller manifest + a 'keep my videos/reels?' prompt.

69 - TELEMETRY / APP-LEVEL OBSERVABILITY for the distributed binary is absent. The repo has per-video emit_indexer_report and the M8 cost ledger, but there is NO crash-reporting, NO 'phone home on failure' (or an explicit decision NOT to, with the privacy rationale), and NO way for a non-technical user to send a diagnostic bundle. The owner's own 'launch quality bar' memo demands first-class observability across every NEW code path (installer/bootstrap/torch-fetch/weight-fetch) ? none of that new code is exercised by the current offline CI, so failures in the field are invisible.

70 - WINDOWS-SPECIFIC HAZARDS are barely touched despite the target persona being a non-tech India coach on Windows. Missing: (a) MAX_PATH/long-path failures when the managed conda env + nested torch/CUDA dirs blow past 260 chars under %LOCALAPPDATA%; (b) CRLF/encoding when the wizard writes config.json/.env; (c) Windows Defender/SmartScreen quarantining an unsigned .exe or PyInstaller onefile (AV false-positives are mentioned only in passing for onefile); (d) the documented port-8000-busy case has a free-port fallback only in the tray-launcher slice, not the bare `indexer --server` path; (e) spaces/non-ASCII in Windows usernames breaking conda/uv paths.

71 - AUTO-UPDATE mechanics are asserted ('update like Ollama: run a new installer over the old') but never designed: how the running tray app detects a new version, whether it stops the server cleanly before replacing files, version-pinning the GCS WASB weights against the new engine, and (critically) reconciling auto-update with the DB-migration gap above. 'Run installer over old' silently assumes the data dir survives ? which contradicts the still-unresolved CWD-vs-app-home question.

72 - FIRST-RUN WIZARD is described as living in the SPA across multiple research topics but the engine has NONE of the endpoints it needs. GET /api/capabilities (the wizard's compute-step source) DOES NOT EXIST ? confirmed by grep: main.py exposes only /api/config, /api/health, /api/process, /api/jobs, /api/query, /api/compile, /api/highlights, uploads. The wizard plan is built on three net-new endpoints (/api/capabilities, /api/videos, /api/reels) that are unbuilt and owner-gated. The plan treats the wizard as 'config wiring, not net-new' ? that is FALSE for the capabilities endpoint it depends on.

73 - SPA<->ENGINE VERSION SYNC has no mechanism. khelsutra/web/package.json version is '0.0.0' and the client hardcodes const BASE='/api' (client.ts:7) with NO API-version header, NO

/api/version handshake, and NO contract test between the bundled SPA and the engine it ships with. The plan flags 'version skew' as a risk but proposes no concrete versioned-contract test, no engine-advertised API version, and no build-time pin recorded in the wheel (dist/ is gitignored, so there is NO auditable record of which SPA shipped in a given engine build).

74 - TRAINING GATING / LONG-RUNNING JOB UX is unsolved for a web app. worker train is an OFFLINE CLI; the in-process FastAPI job worker is a ThreadPoolExecutor whose jobs DO NOT survive a restart (main.py:646 fail_stale_jobs sweeps them to 'failed'). A multi-hour GPU training run cannot run as one of these ephemeral in-process jobs, yet the plan never specifies how training is launched from the web UI, how progress is streamed, how it is cancelled, or how it survives the tray app being quit. The honesty fence (gate_verdict hardcodes ship=False at harness.py:177 ? confirmed) must also be surfaced in-UI, which has no design.

75 - ANNOTATION ('rally annotation if needed', an explicit GOAL) has no in-app path. The only annotator is the external VLC 3.0.x Lua plugin in a sibling repo. The plan acknowledges this but does NOT decide between bundling VLC (heavyweight, version-pinned to 3.0.x) vs building the 'HTML5 <video>' in-browser annotator ? leaving a stated product goal with no implementation path and no owner decision.

76

77 [missing_topics]

78 - Concurrency / single-instance enforcement: nothing prevents two launches (tray double-click + a stale `indexer --server`) from opening the SAME SQLite DB. WAL is on (database.py:41) but the in-process executor + fail_stale_jobs sweep (main.py:646) assume ONE writer ? a second process would sweep the first's live jobs to 'failed'. No PID/lockfile or single-instance guard is planned.

79 - Disk-space / quota management for a long-lived desktop install: output reels, scratch fetches of bucket inputs, and cached weights accumulate. The repo HAS cache-hygiene tooling (tools/cleanup_caches.py, disk guard per the memory index) but the packaging plan never wires it into the app (no UI surface, no auto-purge default, no first-run disk-space check before a GCS weight download).

80 - Weight integrity + supply-chain: the recommendation says 'verify a pinned SHA-256' for the WASB download but no plan exists for WHERE the canonical hash is stored, what happens on mismatch (refuse vs warn), or signing. Combined with WASB-C3 (provenance UNRESOLVED, ship-blocking ? Dockerfile lines 14-16 + harness.py:177 hardcode ship=False), auto-fetching these weights into a PUBLIC download likely VIOLATES the MIT/Apache/BSD-only guardrail until C3 is cleared ? the plan never resolves whether the public app may fetch them at all.

81 - scipy is an UNDECLARED serving dependency (transitive via supervision only) ? flagged in the verification section but never tracked as a concrete pre-packaging task; a freeze that prunes unreferenced transitives or a supervision bump silently breaks inference. Must become an explicit pyproject dep before any freeze/bundle.

82 - Built-artifact CI lane: today's 5 CI lanes test SOURCE on Python 3.13 + CPU torch; NO lane builds and smoke-tests the frozen/bundled artifact (hidden imports, data files like backend/assets/fonts, namespace-package resolution since backend has no __init__.py). Packaging regressions are invisible until an owner manually runs the binary. No plan to add this lane or to exclude bootstrap code from the 70% coverage floor.

83 - Default-segmenter contradiction with offline-first: README config defaults default_segmenter to gemini, but a zero-config local download must default to the truly-local preset (skip_ai_handoff=True). The plan never specifies the SHIPPED config.json / DEPLOYMENT_PROFILE default, so the out-of-box experience could surprise-bill users on Gemini or 0-rally on WASB-without-key.

84 - CORS posture in the packaged single-origin app: main.py:64-67 defaults allow_origins=['*'] with credentials off. For a localhost appliance this is mostly moot, but a DNS-rebinding attack against 127.0.0.1:8000 + the unauthenticated GET /api/config (full config dump) is a real local-exploit vector the plan never addresses (no Host-header check, no localhost-origin allowlist by default).

85 - Conda/Anaconda ToS is NOT yet clean in the shipped reference image: deploy/cloudrun/Dockerfile uses Miniconda3 + `conda tos accept ... pkgs/main`/`pkgs/r` (lines ~33-35), exactly the Anaconda-Distribution-ToS exposure the plan flags. The recommended micromamba+conda-forge fix is NOT implemented anywhere yet and isn't tracked as a blocking task for the WASB env in any artifact.

86 - Time-to-first-result / model-download UX on Indian bandwidth: a multi-hundred-MB weight fetch + torch wheel (CUDA wheels are ~2-2.5GB) on first run needs resumable downloads, a CDN/mirror decision, and a clear progress UI ? mentioned in passing but with no concrete resumability/mirror plan or a size budget per tier.

87 - Localization quality gate: kn/hi catalogs are AI-DRAFTED with native-speaker review PENDING (flagged in the digest). Shipping the trilingual UI to real non-test users before review is a trust risk with no gating decision in the packaging plan.

88 - Backup/restore of the user's index + labelled corpus: a BYO-training app accumulates
irreplaceable human labels (rally CSVs) and a trained-model corpus. No export/import/backup
feature is planned, and the MD5 video<->label join means re-encoding breaks it ? a data-loss
landmine with no safeguard.

89

90 [unverified_or_refuted]

91 - REFUTED ? 'no Dockerfile exists' (still asserted in
khelsutra/docs/LOCAL_APPLIANCE_ARCHITECTURE.md:49,96,135,140 as '[verified]'): a Cloud Run GPU
Dockerfile DOES exist at deploy/cloudrun/Dockerfile (shipped in #287/M6). The prior-art-deepdive
research already caught this, but the live appliance DESIGN DOC still encodes the stale claim in
four places including its DoD ? recommendations built on 'feasibility-check a worker image
before promising Cloud Run' are arguing a settled point.

92 - WEAKENED ? Nuitka is described as 'Apache-2.0' in the frozen-binaries fit narrative
but the refuted note correctly states current Nuitka 4.x is AGPL-3.0+ (with a binary runtime
exception). Any recommendation text that lists Nuitka as 'Apache-2.0 (compatible)' (it appears
that way in the verification-section implications) is wrong and must be corrected to
AGPL-3.0+/with-exception before it informs a licensing decision.

93 - UNVERIFIED ? 'uv pip install --torch-backend=auto into the tool's venv' is presented
as the spine of the lead distribution channel, but it is also stated that --torch-backend=auto
does NOT work in uv tool install/uvx (only uv pip / project mode). The proposed installer script
reaches into the tool's venv to run uv pip ? this cross-mode hack is asserted to work but is not
verified against a real uv version, and uv's tool-env layout is not a stable public contract.
High risk the lead channel's central mechanic doesn't work as written.

94 - UNVERIFIED ? 'WASB-SBDT is MIT, so weights are license-clean' is repeatedly used to
justify shipping/fetching wasb_badminton_best.pth.tar, but the CODE-LEVEL truth
(Dockerfile:14-16, harness.py:175) is that the CHECKPOINT provenance is C3-UNRESOLVED (academic-
data-trained) and ship is hardcoded False. The MIT CODE license does NOT clear the WEIGHTS ?
conflating them weakens every 'fetch weights at first run, license-clean' recommendation. Do not
let MIT-code stand in for weight provenance.

95 - UNVERIFIED ? 'imageio-ffmpeg lacks ffmpeg and is a GPL build' is stated as fact to
reject it; plausible but not grounded in a file/source in this pass. Since it drives the ffmpeg-
delivery recommendation, it should be verified before being treated as decisive.

96 - UNVERIFIED ? the claim that the truly-local native-WASB config 'that produced 22
rallies on an L4' is reproducible is contradicted by the noted prior cold-start config-drift
that yielded 0 rallies. The exact reproducible config/command is called a 'missing piece' yet a
'native' packaged profile is proposed to bake it in. Shipping that profile on an unverified
config risks a 0-rally out-of-box failure.

97 - WEAKENED ? 'pystray is the tray lib' is recommended while simultaneously noting
pystray is LGPL-3.0, which trips the MIT/Apache/BSD-only guardrail. The recommendation hedges
('resolve the license question first') but still sequences a slice around pystray; an LGPL
dependency should be rejected up front per the non-negotiable guardrail, not provisionally
adopted.

98

99 [contradictions]

100 - DISTRIBUTION CHANNEL vs DUAL-ENV REALITY: the distribution-channels topic leads with
'uv tool install' as the primary channel, but the dual-python-env and frozen-binaries topics
both conclude NO single-venv channel can host both torch ABIs (engine torch 2.6+cu124 vs WASB
torch 1.11+cu113). The lead channel therefore only ever delivers the LIGHT tier; the plan
presents uv-first as if it solves distribution when it explicitly cannot deliver the headline
WASB-quality CUJ.

101 - D10 'SELF-HOSTED WEBSERVER, NOT A DESKTOP APP' vs the LAUNCHER+INSTALLER+SYSTEM-TRAY
recommendation: launcher-ux argues a tray app 'satisfies D10's intent', but a Start-Menu
shortcut + tray icon + auto-open-browser + Inno Setup installer IS a desktop application by any
normal-user definition. The plan relitigates a superseded decision (DESKTOP_APP_PLAN) under a
euphemism instead of getting an explicit owner re-ratification.

102 - BIND ADDRESS: the containerization topic CORRECTLY states the server must bind
0.0.0.0 inside a container (confirmed: main.py:756 hardcodes uvicorn.run(host='127.0.0.1')),
while the cloud-one-click and completed-CUJs topics treat the 127.0.0.1 bind as a SECURITY
INVARIANT that must be preserved. These are not reconciled ? the same hardcoded line is
simultaneously 'a required change' and 'a must-keep default,' with no config knob proposed to
make bind address a setting gated by the edge-auth profile.

103 - 'BATTERIES-INCLUDED' vs CLOUD-SERVING SELF-CONTAINMENT: the config topic states
cloud-serving in a downloaded app is NOT self-contained (needs the user's own GCP project + a
deployed Cloud Run image + ADC), yet the cloud-one-click topic markets a 'push my workload to
cloud' toggle and D13 'your machine or our cloud, switch anytime.' Either there is a hosted

Khelsutra service (no StorageBackend/ComputeBackend client exists for it today) or the cloud half is BYO-GCP ? the plan markets (a) while only (b) is built.

104 - WEIGHTS LICENSE-CLEAN vs SHIP-BLOCKED: weights-and-ffmpeg-bundling recommends auto-fetching WASB weights at first run as 'license-clean (MIT)', while verification and the prior-art topic both state WASB-C3 provenance is UNRESOLVED and gate_verdict hardcodes ship=False for any non-owner serving. A PUBLIC download that auto-fetches those weights directly contradicts the non-negotiable weights-licensing guardrail; the plan cannot both 'fetch at first run' and 'respect C3' until C3 is cleared.

105 - CPU-DEFAULT vs PRESERVE-EVERY-CUJ: multiple topics recommend shipping a CPU/Gemini default and gating WASB behind an opt-in heavy tier, but the stated GOAL is 'preserve every already-completed end-to-end CUJ' ? and the verified best-quality CUJ (CUJ-6/7) IS native WASB on GPU. A CPU-only default bundle SILENTLY DROPS a completed CUJ; the plan does not square 'preserve every CUJ' with 'default has no WASB.'

106

107 [followups]

108 - DECIDE THE APP-DATA ROOT NOW and write the one-PR change: introduce a resolved app-home (env override + OS default), and thread it through loader.DEFAULT_CONFIG_PATH, the load_dotenv(path=...) call at main.py:10, OutputConfig.output_dir, the db path, and the static mount at main.py:71-72. Unify doctor.py's REPO_ROOT config write (scripts/doctor.py:54) with the server's read so --setup and the server agree. This is the prerequisite for uninstall, auto-update, and migration ? sequence it first.

109 - Make the bind address and CORS origin profile-driven config (not hardcoded): replace uvicorn.run(host='127.0.0.1') at main.py:756 with a config value defaulting to 127.0.0.1, only widening to 0.0.0.0 when an edge-auth profile is active, and add a Host-header / localhost-origin guard so the unauthenticated GET /api/config dump (main.py:144) is not reachable via DNS-rebinding.

110 - AUDIT GET /api/config for secret leakage and add field redaction NOW, before any packaging ? it returns the entire AppConfig unauthenticated. Add a test asserting no credential-shaped field is ever serialized, so a future config field can't silently leak the Gemini key.

111 - Resolve WASB-C3 (or formally scope it OUT of v1): get counsel sign-off on whether the checkpoint may be redistributed/auto-fetched into a PUBLIC build. Until resolved, the only honest v1 is LIGHT-tier (Gemini/motion/stub) ? make that the explicit, recorded v1 scope and ship the WASB path only to the owner. gate_verdict (harness.py:177) and the Dockerfile note already encode ship=False; the packaging plan must inherit that fence.

112 - Build the three net-new endpoints the wizard/console depend on as a single tracked engine slice BEFORE the packaging push: GET /api/capabilities (honest hardware/segmenter probe), GET /api/videos (list, today only single-row get_video), GET /api/reels/{video_id} (reload-survival). Confirmed absent in main.py. Add an /api/version handshake in the same slice and a contract test pinning the bundled SPA against it.

113 - Replace Miniconda + `conda tos accept` in deploy/cloudrun/Dockerfile (lines ~33-35) with micromamba + conda-forge for the WASB env, and record the decision in an ADR; this is the license-clean path for BOTH the cloud image and any local managed env, and it's currently unimplemented.

114 - Add scipy as an explicit [project.dependencies] entry in pyproject.toml before any freeze/bundle (currently transitive via supervision only) and add an automated license-scan CI lane (pip-licenses/reuse) given the guardrail is human-only today ? a single PR closes both latent gaps.

115 - Add a built-artifact smoke CI lane: build the chosen artifact (uv tool install / onedir freeze / container), then run `indexer --help` and boot the server + hit /api/health, so packaging regressions (hidden imports, namespace-package resolution with no backend/__init__.py, missing data files) are caught in CI rather than by the owner manually.

116 - Define the DB migration + backup contract for upgrades: on app start, back up the SQLite file before running the user_version ladder (database.py), refuse to open a DB whose user_version exceeds the binary's known max (downgrade guard), and add an explicit 'export my data' affordance. Pair this with the auto-update design so 'run installer over old' is provably safe.

117 - Get an explicit owner re-ratification of D10 vs the tray-launcher recommendation in writing, and resolve the pystray LGPL question by choosing an MIT/BSD tray library (or stdlib webbrowser-only) up front rather than provisionally ? the guardrail forbids LGPL deps, full stop.

118 - Write a single 'PACKAGING_PLAN.md' tracked project doc (per the repo's tracked-project-doc convention) that fixes the v1 SCOPE decisions: which tier is v1 (recommend LIGHT/Gemini), shipped default profile/config.json, whether training + annotation are in v1 (recommend deferred given ship=False + VLC dependency), and the explicit out-of-scope list ? so the contradictions above are decided once, not relitigated per channel.

28. assistant (2026-06-21T14:13:13.370Z)

The workflow gave me a deeply-grounded synthesis plus a sharp critic. Before I commit these to a plan doc, the project's own lesson ([[lesson-verify-engine-surface-before-scoping]]) says I must re-verify load-bearing engine claims at `file:line` myself ? the engine ships PRs between sessions. Let me verify the P0-critical claims directly.

29. user (2026-06-21T14:13:17.661Z)

```

=== /api/config handler ===
144:@app.get("/api/config")
145:def get_config():
146:    return global_config

=== uvicorn bind ===
756:    uvicorn.run(app, host="127.0.0.1", port=args.port)

=== net-new endpoints presence (should be ABSENT) ===
347:@app.get("/api/videos/{video_id}/cost")

=== existing api routes ===
74:@app.get("/")
144:@app.get("/api/config")
148:@app.get("/api/health")
281:@app.post("/api/process")
332:@app.get("/api/jobs/{job_id}/cost")
347:@app.get("/api/videos/{video_id}/cost")
358:@app.get("/api/jobs/{job_id}")
391:@app.get("/api/highlights/{video_id}/{filename}")
415:@app.post("/api/query")
426:@app.post("/api/compile")

=== load_dotenv + config path + static mount ===
backend/main.py:8:from dotenv import load_dotenv
backend/main.py:10:load_dotenv()
backend/main.py:24:from fastapi.staticfiles import StaticFiles
backend/main.py:71:os.makedirs("backend/api/static", exist_ok=True)
backend/main.py:72:app.mount("/static", StaticFiles(directory="backend/api/static"),
name="static")
backend/main.py:76:    index_path = "backend/api/static/index.html"
backend/main.py:80:    return "<h3>Sports Highlight Indexer Server is Running. Please create
index.html in backend/api/static/</h3>"
backend/config/loader.py:30:DEFAULT_CONFIG_PATH = Path("config.json")
backend/config/loader.py:112:    cfg_path = Path(path) if path is not None else
DEFAULT_CONFIG_PATH
backend/config/loader.py:202:    cfg_path = Path(path) if path is not None else
DEFAULT_CONFIG_PATH

```

30. assistant (2026-06-21T14:13:27.041Z)

Core P0 claims confirmed exactly as the workflow reported. Let me verify the licensing/dual-env/ship-gate claims that underpin the tier decisions.

31. user (2026-06-21T14:13:32.746Z)

```

=== scipy usage (claimed transitive, undeclared) ===
backend/eval/rally_seq_proto.py:35:from scipy.ndimage import uniform_filter1d
backend/eval/rally_seq_proto.py:36:from scipy.stats import rankdata
backend/eval/tolerance_metrics.py:68:    from scipy.optimize import linear_sum_assignment
backend/features/rally_seq.py:17:from scipy.ndimage import uniform_filter1d, maximum_filter1d

=== WASB shell-out to $WASB_PYTHON ===

```

```

17:    (`WASB_WEIGHTS` / `WASB_REPO_DIR` / `WASB_PYTHON` / `WASB_SCRATCH`), then
33:(a real GPU + the `wasb` env), so the offline, subprocess-mocked suite asserts the
contract,
39:import subprocess
69:    weights_path: str = ""                                # REQUIRED (env WASB_WEIGHTS or
indexing.wasb.weights_path)
89:        python_bin=env("WASB_PYTHON") or w.get("python_bin", "python"),
90:        weights_path=env("WASB_WEIGHTS") or w.get("weights_path", "") or "",
110:    def _run(self, argv: List[str], cwd: Optional[str] = None) ->
subprocess.CompletedProcess:
112:        return subprocess.run(argv, cwd=cwd, capture_output=True, text=True,
117:        """Probe `torch.cuda.is_available()` in the wasb env's python (a subprocess). Any
failure
121:        except (FileNotFoundError, subprocess.TimeoutExpired):

=== ffmpeg encoder call sites (libx264) ===
39:    """Make a filesystem path safe to embed inside single quotes in an ffmpeg
43:    The colon MUST be escaped even inside single quotes: ffmpeg's filtergraph
126:    """Build the ffmpeg -vf filtergraph that burns the configured watermark (text
203:    success the clip path and None; on a skip (invalid/out-of-range/ffmpeg failure)
249:        "ffmpeg", "-y",
253:        "-c:v", "libx264",
279:        f"ffmpeg: {_err or '(no stderr captured)'}"
288:        last_exc = None # ffmpeg succeeded but produced nothing
298:        return None, (idx, start, end, f"ffmpeg error: {stderr_msg[-200:]}")
369:        f"Common causes: video file is corrupted, segments are out of range, or
ffmpeg is not installed."

=== Anaconda ToS in cloudrun Dockerfile ===
11:# * a conda `wasb` env (py3.8 + torch 1.11.0+cu113) runs the WASB shuttle detector as a
subprocess
23:    CONDA_DIR=/opt/conda \
33:RUN curl -sSL https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh -o
/tmp/mc.sh \
35:    && "$CONDA_DIR/bin/conda" tos accept --override-channels --channel
https://repo.anaconda.com/pkgs/main || true \
36:    && "$CONDA_DIR/bin/conda" tos accept --override-channels --channel
https://repo.anaconda.com/pkgs/r || true \
37:    && "$CONDA_DIR/bin/conda" create -y -q -n wasb python=3.8 \
52:# The native runner finds the WASB env + repo by env (no `conda activate` needed); weights
come in at
54:ENV WASB_PYTHON=/opt/conda/envs/wasb/bin/python \

=== gate_verdict ship=False ===
1:"""Gen-0 TAL harness (M0.4): ONE command -> train -> honest LOVO eval -> ship-gate ->
artifact.
17: (WASB C3 unresolved => commercial ship refused; explicit F1@tIoU target still
142:def gate_verdict(eval_res: dict, floor: Optional[dict]) -> dict:
143:    """Pre-flight ship-gate.
151:    ship=True only when all non-C3 blocks clear (honesty first).
174:    reasons.append("ship-gate F1@tIoU target NOT YET CALIBRATED ? provisional; see issue
#172")
175:    reasons.append("WASB checkpoint provenance C3 UNRESOLVED ? commercial ship blocked
regardless of F1")
176:    # ship remains False until C3 resolved + other gates (by design).
177:    return {"ship": False, "floor_passed": floor_passed, "sign_test": sign, "blockers":
reasons}
196:    verdict = gate_verdict(eval_res, floor)

=== DB user_version ladder ===
69:    `user_version` stays un-bumped ? and the next open would re-run the same
88:    # Versioned migrations live here. PRAGMA user_version is the schema
91:    version = cursor.execute("PRAGMA user_version").fetchone()[0]
94:    self._migrate_to_v1(cursor)
97:    self._migrate_to_v2(cursor)

```

```

100:         self._migrate_to_v3(cursor)
103:         self._migrate_to_v4(cursor)
106:         self._migrate_to_v5(cursor)
109:         self._migrate_to_v6(cursor)
110:         # future: if version < 7: cursor.execute("ALTER TABLE ..."); PRAGMA user_version
= 7
155:     def _migrate_to_v1(self, cursor):
183:         cursor.execute("PRAGMA user_version = 1")
185:     def _migrate_to_v2(self, cursor):
210:         cursor.execute("PRAGMA user_version = 2")
212:     def _migrate_to_v3(self, cursor):
222:         cursor.execute("PRAGMA user_version = 3")
224:     def _migrate_to_v4(self, cursor):
242:         cursor.execute("PRAGMA user_version = 4")
244:     def _migrate_to_v5(self, cursor):
280:         CREATE UNIQUE INDEX IF NOT EXISTS idx_user_models_user_version

```

32. assistant (2026-06-21T14:13:57.303Z)

Every load-bearing claim verified at `file:line` ? the synthesis is accurate:

- `/api/config` returns raw `global\_config` unauthenticated (main.py:144-146); uvicorn hardcodes `127.0.0.1` (main.py:756); `/api/capabilities`, `/api/videos` (list), `/api/reels` are **absent**; all state is CWD-relative (loader.py:30, main.py:10/71-72).
- scipy is imported on the **serving** path (`backend/features/rally\_seq.py:17`) but **undeclared** in pyproject/requirements.
- WASB shells out to `\$WASB\_PYTHON` (native_wasb_runner.py:89); encoder is GPL `libx264` (compiler.py:253); the Cloud Run image uses Miniconda + `conda tos accept` (Dockerfile:33-37); `gate\_verdict` hardcodes `ship=False` on WASB-C3 (harness.py:177); DB is at `user\_version=6`.

Let me check the project doc template + the superseded plan's header so the tracked doc follows convention and records lineage.

33. user (2026-06-21T14:13:58.117Z)

```

1      <!--
2      PROJECT DOC TEMPLATE ? copy this file to docs/<YOUR_PLAN>.md and fill it in.
3
4      A "tracked project doc" is how this repo carries substantial design/project work: it can
be
5      ITERATED on, its progress TRACKED, and it is MARKED DONE against explicit deliverables,
then
6      archived. Standardizes the emergent pattern in docs/RALLY_QUALITY_RESEARCH.md §7 and
7      docs/GOLDEN_REGRESSION_FIXTURES.md.
8
9      Rules of thumb:
10     - §7 (the progress tracker) is the SOURCE OF TRUTH ? one row per independently-shippable
PR.
11     - Keep it HONEST: tag each claim [verified] (checked against code), [researched] (needs
sign-off),
12     or [design] (proposed). The doc reflects real shipped state, never aspirational.
13     - It POINTS to detail (code anchors, research artifacts), it does not duplicate it.
14     - Sections marked OPTIONAL are included only when relevant (e.g. Foundation/Threat-model
for
15     research- or security/privacy-heavy work). Delete sections that do not apply.
16     - Move to docs/archives/ once Status = DONE (per the archive policy: only terminal-state
work).
17     -->
18
19     # <Project / Plan Title>
20
21     > **Status:** `DRAFT` ? `<owner-gated? | in progress>` . **Owner:** `<name>` .
**Created:** `<YYYY-MM-DD>` . **Last updated:** `<YYYY-MM-DD>`
22     > **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when
DONE)
23     > **Tracking anchors:** §7 progress tracker is the source of truth; indexed in

```

```

`docs/README.md`; pointer in memory `current-state` (and `docs/NEXT_STEPS.md` once work starts).
24 > **Relation to existing docs:** `<extends | supersedes | peer-of ?>`
25 > **Honesty note:** mark each claim `[verified]` / `[researched]` / `[design]`. Not
legal/financial advice where applicable.
26
27 ---
28
29 ## 0. TL;DR
30 <2?5 sentences: the goal, the approach, and the single most important caveat. Call out
any finding that is already actionable on its own.>
31
32 ## 1. Problem & goal
33 <Why now; the concrete user-facing outcome; what "good" looks like. Link the code/docs a
reader must read to get current.>
34
35 ## 2. Decisions locked
36 <Table ? capture answered questions so they are not relitigated.>
37
38 | # | Decision | Source / date | Implication |
39 |---|-----|-----|-----|
40 | D1 | ? | ? | ? |
41
42 ## 3. Foundation ? research / prior art *(OPTIONAL)*
43 <Include when the plan rests on external facts: cited research, legal/compliance,
benchmarks, prior-art survey. Mark researched-vs-verified; list any sign-off gates this
creates.>
44
45 ## 4. Design / architecture
46 <The actual plan. Data contracts, diagrams, and a REUSE MAP ? name the existing
modules/code this builds on so it is clear what is new vs reused.>
47
48 ## 5. Threat model / risk table *(OPTIONAL ? security / privacy / compliance work)*
49 <Who/what is protected vs NOT; residual risks stated plainly.>
50
51 ## 6. Honest limits ? what this does NOT do
52 <The false-confidence guardrail: what a green suite / a DONE status does NOT prove.>
53
54 ## 7. Deliverables & progress tracker ? **source of truth**
55 Legend: ? Todo . ? In progress . ? Done . ? Blocked/gated. **One small PR per row**,
branched from `origin/master`, no stacking. Update the row (status + PR link) as work lands.
56
57 | ID | Deliverable | Depends on | Gated? | Status | PR |
58 |---|-----|-----|-----|-----|----|
59 | P0 | ? | ? | No | ? | ? |
60
61 ## 8. Open questions ? owner / external *(split blocking-gates from nice-to-knows; name
who decides)*
62 1. ?
63
64 ## 9. Definition of done
65 <Explicit acceptance checklist. When all ticked ? flip `Status ? DONE` and archive.>
66 - [ ] ?
67
68 ## 10. References
69 **Internal code anchors `[verified]`:** `<paths>`. **Internal docs:** `<docs>`.
**External:** `<cited sources>`.
70
71 ---
72
73 ### Changelog
74 - `<YYYY-MM-DD>` ? `<what changed>`.
75

```

All verified. Now I'll write the tracked project doc (per the repo's convention ? the workflow's own top followup) and then give you the plan in chat. Writing it as an untracked file; I won't commit or open a PR (owner-gated).

35. assistant (2026-06-21T14:16:18.189Z)

Batteries-Included App ? Packaging & Distribution Plan

```
> **Status:** `DRAFT` ? owner-gated (scope call 02 is load-bearing) . **Owner:** Avi Dullu .
**Created:** 2026-06-21 . **Last updated:** 2026-06-21
> **Lifecycle:** `DRAFT` ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when DONE)
> **Tracking anchors:** $7 progress tracker is the source of truth; index in `docs/README.md`;
pointer in memory `current-state` + `docs/NEXT_STEPS.md` once work starts.
> **Relation to existing docs:** **supersedes** the narrative of `docs/DESKTOP_APP_PLAN.md` (per
locked D10 ? self-hosted webserver, not a desktop app); **builds on**
`docs/COMPUTE_DECOUPLED_SERVING/` (the profile/storage/compute seam) and is the **distribution
layer** under `kheIsutra/docs/LOCAL_APPLIANCE_ARCHITECTURE.md`.
> **Honesty note:** claims tagged `[verified]` (checked at file:line this session),
`[researched]` (web/needs sign-off), `[design]` (proposed). Not legal advice ? licensing items
need counsel.
```

0. TL;DR

Ship the engine as ****ONE pip-installable artifact** whose behaviour is selected entirely by the already-merged decoupled-compute seam¹ (``deployment.profile` / `compute_target` / Storage+Compute backends`), fronted by a friendly launcher + a Windows installer ? **\*\*not\*\*** a frozen monolith, **\*\*not\*\*** an Electron/Tauri shell. The hard constraint that forces scope: torch is deliberately out of ``pyproject`` (needs ``--index-url``) **\*\*and\*\*** native WASB needs a *\*second, incompatible\** torch env (conda py3.8 + torch 1.11+cu113) shelled out via ``$WASB_PYTHON`` ``[verified native_wasb_runner.py:89]`` ? so no single venv/freeze can deliver the GPU-quality path. ****The honest v1 is a LIGHT tier**** (Gemini + motion + cpu-mock, **zero** torch/GPU) that preserves CUJ-1.1.5/8; the native-WASB GPU path (CUJ-6/7) ships ****owner-only**** until two blockers clear (WASB-C3 weight provenance + a reproducible truly-local config). ****Five small engine prerequisites must land before any packaging push**** (\$7 P0). The single most actionable finding: ****all mutable state is CWD-relative today and ``--setup``/server disagree on ``config.json`` ? a silent data-loss footgun that P0 must fix first.**

1. Problem & goal

****Goal:**** a non-technical user (the primary persona: an India coach/parent on Windows ? ``docs/UI_PROPOSAL.md`` \$2) ****downloads and runs**** the indexer ****without sourcing the repo****. The app spins up a ****local web service**** for (a) inference video?rallies?reel, (b) BYO ****training****, and (c) ****rally annotation**** if needed ? runnable ****locally and against cloud services**** ? and ****preserves every completed CUJ**** (\$ CUJ map).

****Why now:**** scaffolding + the decoupled-compute serving seam (M1?M8) are merged and default-OFF; the SPA + i18n + the dry-run CUJ are done. What's missing is the **distribution** story: today the only path is ``git clone`` ? ``pip install -e .`` ? out-of-band torch ? manual ffmpeg ? set up the WASB conda env ? build the SPA from a **different repo**. This plan collapses that into download-and-run.

```
**Read to get current:** `pyproject.toml`, `backend/main.py` (`main()` + `--server`),
`backend/config/{loader,presets,models}.py`, `backend/storage/*`, `backend/compute/*`,
`backend/pipeline/detectors/native_wasb_runner.py`, `deploy/cloudrun/Dockerfile`,
`docs/COMPUTE_DECOUPLED_SERVING/README.md`, `kheIsutra/web/` +
`kheIsutra/docs/LOCAL_APPLIANCE_ARCHITECTURE.md`.
```

2. Decisions locked (do not relitigate)

#	Decision	Source / date	Implication for packaging
L1	**Pure DETECT?WINDOW?STITCH core never branches on profile** ; one build, profile-selected		
	COMPUTE_DECOUPLED_SERVING D1	Local vs cloud is **one config flip** , not a fork ? build ON	

this, don't reinvent. |
 | L2 | **Profiles** = truly-local / cloud-serving / cpu-mock via `PROFILE\_PRESETS` | M1 #279
 `[verified backend/config/presets.py]` | The app exposes a single "where does compute run"
 choice. |
 | L3 | **Compute** = GCP only** (Cloud Run GPU Jobs); DigitalOcean dropped | ADR-013, 2026-06-20 |
 Cloud tier targets Cloud Run + BYO-GPU-VM only. |
 | L4 | **Edge-auth** required before any non-localhost exposure** (Cloudflare Access + in-app JWT
 verify) | D9 / M9 `[verified backend/api/auth.py wired default-OFF main.py:294]` | Bind stays
 `127.0.0.1` by default; `0.0.0.0` only under an edge-auth profile. |
 | L5 | **Delivery** = self-hosted webserver the user spins up ? NOT a desktop app** | D10,
 2026-06-20 (supersedes `DESKTOP\_APP\_PLAN.md`) | Launcher opens the user's *own browser* at
 localhost; see 01 (re-ratify vs installer/launcher). |
 | L6 | **MIT / Apache-2.0 / BSD only** for every dependency ? code, data, AND weights | ratified
 2026-06-09 | Every *new packaging* dep must clear this (see S3). |
 | L7 | **Paid LLMs (Gemini)** = inference/serving only, never a training data source** ; commercial
 models trained on paid-LLM labels are ship-blocked | guardrail + `gate\_verdict` `[verified
 harness.py:177 ship=False]` | The wizard's Gemini-key step must say so; training UI must surface
 the ship-gate. |
 | L8 | **User-switchable profiles** ("your machine or our cloud") with explicit confirm ? cloud
 spend never silent | D13 / D15 | "Push to cloud" toggle needs confirm + INR estimate from
 `backend/billing.py`. |

3. Foundation ? research / prior art

Distribution channels `[researched]`** ? `uv` is the only channel that can both install a
 pinned Python *and* route torch's `--index-url`; `pipx`/plain `pip` can't express the torch
 index. Frozen binaries (PyInstaller/Nuitka) with torch+CUDA are 3?5 GB, can't split CUDA libs,
 and **cannot host the second WASB env** ? rejected for the GPU path. Prior art (Frigate NVR ?
 closest analog: ffmpeg + ML + web + Docker; Ollama; ComfyUI; LM Studio) converges on **two**
 patterns that fit us:** (1) a **container image** for the server/GPU profile, (2) a **first-run**
 download + a friendly launcher** for the local app.

Licensing landmines `[verified]` (the load-bearing S3 findings):**
 - **WASB-C3** is the dominant blocker.** The WASB-SBDT *code* is MIT, but the badminton
checkpoint provenance is **unresolved** for commercial redistribution ? `gate\_verdict`
 hardcodes `ship=False` `[verified training/gen0/harness.py:174-177]`. **Do not conflate MIT-code**
 with cleared-weights.** A public download must **not** auto-fetch the weights until counsel
 signs off (03).
 - **ffmpeg** is GPL via the `libx264` encoder** `[verified
 backend/pipeline/stitcher/compiler.py:253]`. The license-clean path is **not** to redistribute
 ffmpeg ? resolve a system/winget install instead. `imageio-ffmpeg` also lacks `ffprobe` (which
 the engine needs) `[researched]`.
 - **Anaconda ToS exposure**: the Cloud Run image uses `Miniconda3` + `conda tos accept` ?
 pkgs/main`/`pkgs/r` `[verified deploy/cloudrun/Dockerfile:33-37]` ? the 200-employee ToS trap.
 Replace with **micromamba + conda-forge** (BSD, no ToS) for the cloud image *and* any local
 managed WASB env; record an ADR.
 - **New packaging** deps that clear L6 `[researched]`:** `uv` (Apache/MIT), Inno Setup
 (permissive), `onnxruntime` (MIT), `micromamba` (BSD-3), Docker/OCI (Apache-2.0), boto3/google-
 cloud-* (Apache-2.0), stdlib `webbrowser`. **Forbidden**: `pystray` (LGPL-3.0). **Build-time-**
only: `Nuitka` 4.x is **AGPL-3.0+** (compiled output has a runtime exception) ? usable as a
 tool, never bundled. `obsreport` (private git dep) stays a truly-optional extra, never bundled.

4. Design / architecture

4.1 One build, profile-selected (the reuse map)

- **Reuse**: hatchling build + `indexer` = backend.main:main` console script `[verified
 pyproject.toml:80-82]`; `PROFILE_PRESETS` `[verified presets.py]`; `get_compute_backend(config)`
 returns `CloudRunCompute` only for `compute\_target=='cloud\_run'` else `None`=in-process
 `[verified backend/compute/\_\_init\_\_.py]`; `backend/storage/\*` (local/gcs/s3/gdrive); the in-
 process `ThreadPoolExecutor(max\_workers=1)` job worker `[verified backend/api/job\_worker.py]` ?
server and worker are the same process for serving `[verified main.py:729-736]`, so a
 packaged launcher just runs `indexer --server`.
 - **New**: the app defaults to **truly-local**; the user toggles profile in Settings (rendered
 from a new `/api/capabilities`) with explicit confirm (L8) + INR estimate before any cloud job.

4.2 How one build runs local + cloud

The local FastAPI server is the **control plane** in every profile. ``truly-local`` runs detect/stitch in-process. ``cloud-serving`` flips ``compute_target=cloud_run`` ? the local control plane dispatches a **Cloud Run GPU Job** (the *same* image, forced inline) and bucket-polls a done-marker; ``storage=gcs`` routes I/O. ``cpu-mock`` gives a \$0 deterministic self-test on first run.

4.3 How the UI is bundled (collapses the cross-repo + reverse-proxy problem)

Build ``khelsutra/web`` (``npm ci && npm run build``) **at package time** and vendor ``web/dist/`` into the wheel via a **hatchling force-include** ? a package dir (e.g. ``backend/ui/``), served by FastAPI with an **SPA history-fallback** route. The SPA already calls **same-origin** ``BASE='/api'`` ``[verified khelsutra/web client.ts:7]`` ? **zero SPA code change** for single-origin serving. Retire the legacy CWD-relative ``backend/api/static`` mount ``[verified main.py:71-72]``; resolve paths via ``importlib.resources``. ``dist/`` is gitignored ? the build must record a **bundle hash** + an ``/api/version`` handshake + a contract test pinning SPA?engine.

4.4 Dual-torch / weights / ffmpeg ? concretely

- **Dual-torch:** do **not** host both ABIs in one venv or freeze. v1 LIGHT carries **no torch** at all. The GPU/WASB HEAVY tier is a **separately-provisioned second env** ? short term via **micromamba + conda-forge** reached by ``$WASB_PYTHON`` exactly as ``native_wasb_runner.py:89`` already does; **strategic fix = export WASB to ONNX** (the ``DetectorRunner`` ABC already names an ``ONNXRunner`` writing the identical ``Frame,Visibility,X,Y`` CSV) so the WASB path needs **neither** torch 1.11 nor 2.x ? ``onnxruntime`` (MIT, tens of MB, CPU+CUDA+DirectML+CoreML) dissolves the dual-env. Owner-gated on a real-GPU byte-parity re-validation.
- **Weights:** never bake ``wasb_badminton_best.pth.tar`` into the artifact (size + C3). Fetch on first run, **SHA-256-pinned + resumable**, into the app-data cache, setting ``$WASB_WEIGHTS``. **Until C3 clears**, the public download must NOT auto-fetch ? owner-only flag.
- **ffmpeg:** route **every** ``ffmpeg/`` ``ffprobe`` call through **one resolver** (config path ? env ? ``shutil.which`` ? fail-loud); first-run setup (extend ``scripts/doctor.py``, which already prints the winget command) detects-and-offers ``winget install Gyan.FFmpeg``. An LGPL ffmpeg bundle (BtbN,openh264/NVENC) is a documented fallback that also requires swapping the ``libx264`` default ? tracked, not v1.

4.5 Tiers

Tier	Audience	Contents	Install UX
----- ----- ----- -----			
LIGHT (v1 default)	Non-tech India coach/parent	on Windows; any CPU box	wheel + bundled SPA + Gemini/motion/cpu-mock; no torch/GPU/weights ; ffmpeg via winget-offer; defaults truly-local; user-supplied Gemini key
Inno Setup .exe	(Start-Menu ? friendly launcher ? auto-opens <code>`127.0.0.1:8000`</code>)	primary; uv generated script cross-platform. First run: cpu-mock self-test ? 3-step wizard.	
HEAVY / local-GPU	(owner + opt-in power user)	Owner w/ golden corpus + CUDA box	LIGHT + micromamba WASB env (py3.8/torch1.11) via <code>`\$WASB_PYTHON`</code> ; main torch out-of-band via uv; weights SHA-256-pinned. Gated on C3 + reproducible config ; ONNX supersedes this env
opt-in "Enable local GPU" step; resumable progress (multi-GB on Indian bandwidth); post-install <code>`torch.cuda.is_available()`</code> validate; falls back to LIGHT.			
CLOUD / server	GCP deploy + BYO-GPU-VM	the <code>`deploy/cloudrun/Dockerfile`</code> image (micromamba-ified); cloud-serving + gcs; behind Cloudflare Access	<code>`gcloud run jobs create`</code> per <code>`deploy/cloudrun/README.md`</code> , or <code>`deploy/gcp`</code> VM scripts. Bind? <code>`0.0.0.0`</code> (config-gated) + edge-auth mandatory.

5. Threat model / risk table (local appliance)

Asset	Protected?	Residual risk	control
----- ----- ----- -----			
Gemini API key	Partially	Today <code>`GET /api/config`</code> returns the whole <code>`AppConfig`</code> unauthenticated <code>`[verified main.py:144-146]`</code> . P0 redacts + adds a serialization test + a Host-header/localhost-origin guard (DNS-rebinding vector against <code>`127.0.0.1`</code>). Key lives in app-home <code>`env`</code> (or OS keychain), never in config; scrubbed on uninstall.	
User videos + human rally labels	Yes (local)	MD5 video?label join breaks on re-encode ? never re-encode after annotating; add export/backup (P2d).	
Cloud spend	Yes	Profile flip needs explicit confirm + INR estimate + per-video cap (L8); never silent.	
SQLite index	Yes	DB migration on upgrade: backup-before-migrate + downgrade guard (P2d).	

6. Honest limits ? what this does NOT do

- v1 (LIGHT) ****does not run native WASB**** ? it ****deliberately drops CUJ-6/7 from the default****. This is the load-bearing scope honesty (02). Quality on a CPU/Gemini box ? the GPU-WASB quality.
- A green offline suite does ****not**** prove the ***built artifact*** works (hidden imports, namespace-package resolution with no ``backend/___init__.py``, missing data files like fonts) ? P0 adds a built-artifact smoke lane.
- "Push to cloud" is ****BYO-GCP**** today (no hosted-Khelsutra Storage/Compute client exists) ? D13 marketing currently overstates a hosted option (04).
- Training is ****not**** a closed loop: ``worker train`` is an offline CLI, ``gate_verdict`` refuses commercial ship, and in-process jobs don't survive restart ``[verified main.py:646 fail_stale_jobs]``. Never advertise `annotate?train` as done.

7. Deliverables & progress tracker ? source of truth

Legend: ? Todo · ? In progress · ? Done · ? Gated. ****One small PR per row****, branched from ``origin/master``, no stacking, default-OFF.

ID	Deliverable	Depends on	Gated?	Status	PR
P0a	**App-data HOME** (<code>`RALLY_APP_HOME`</code> env + OS default <code>`%APPDATA%/Khelsutra`</code> \ <code>`~/local/share/khelsutra`</code>) threaded through <code>`loader.DEFAULT_CONFIG_PATH`</code> , <code>`load_dotenv(path=)`</code> (<code>`main.py:10`</code>), <code>`OutputConfig.output_dir`</code> /db path, the static mount; **unify `doctor.py` REPO_ROOT write with the server read**		No	?	?
P0b	**Redact `GET /api/config`** + test asserting no credential-shaped field serializes + Host-header/localhost-origin guard		No	?	?
P0c	**Config/profile-driven bind + CORS** (default <code>`127.0.0.1`</code> ; <code>`0.0.0.0`</code> only under edge-auth profile) ? replace hardcoded <code>`main.py:756`</code>		No	?	?
P0d	**Declare `scipy`** in <code>`[project.dependencies]`</code> (serving use <code>`[verified features/rally_seq.py:17]`</code> , undeclared) + **license-scan CI lane** (pip-licenses/reuse) + **built-artifact smoke lane** (build ? <code>`indexer --help`</code> ? boot ? <code>`/api/health`</code>)		No	?	?
P0e	**Single-instance guard** (PID/lockfile) so a 2nd launch can't open the same SQLite DB and have <code>`fail_stale_jobs`</code> sweep the live process's jobs	P0a	No	?	?
P1a	**Bundle the SPA single-origin** : build <code>`khelsutra/web`</code> at package time ? vendor <code>`web/dist/`</code> into the wheel (hatchling force-include ? <code>`backend/ui/`</code>) + SPA history-fallback; retire legacy static mount; record bundle hash	P0a	No	?	?
P1b	**Net-new endpoints** the wizard/console need: <code>`GET /api/capabilities`</code> (honest hw/segmenter/ffmpeg/weights probe), <code>`GET /api/videos`</code> (list), <code>`GET /api/reels/{video_id}`</code> + <code>`/api/version`</code> handshake + SPA?engine contract test	P1a	No	?	?
P2a	**ffmpeg resolver helper** wrapping all call sites (<code>`compiler.py`</code> , <code>`video.py`</code>) + first-run cpu-mock self-test	P0a	No	?	?
P2b	**LIGHT installer + launcher** : Inno Setup (uv-managed venv + shortcut ? health-poll launcher ? <code>`webbrowser.open`</code> ; free-port fallback; graceful shutdown) + the uv generated script; winget-offer ffmpeg; ship <code>`config.json`</code> defaulting truly-local; code-signing cert	P2a, P1a			
P2c	**In-UI 3-step first-run wizard** (Gemini key / Compute / Videos), rendered from <code>`/api/capabilities`</code> , en/kn/hi	P1b	No	?	?
P2d	**DB upgrade contract** (backup-before-migrate, downgrade guard, "export my data") + **uninstall manifest** (keep-my-videos? prompt; scrub <code>`.env`</code> secret)	P0a	No	?	?
P3a	**micromamba + conda-forge** replaces Miniconda+conda tos accept in <code>`deploy/cloudrun/Dockerfile`</code> (+ ADR); reuse for the local managed WASB env		No	?	?
P3b	**WASB weight fetcher** (SHA-256-pinned, resumable) ? <code>`\$WASB_WEIGHTS`</code> ? **owner-only flag; public must NOT fetch until C3**	P3a	?	O3 (counsel)	?
P3c	**"Push to cloud" toggle** (flip cloud-serving + confirm + INR estimate + per-video cost cap) + wire worker usage-emission into the M8 ledger	P1b	Owner	(cloud spend)	?
P3d	**M9 edge-auth verified path** (JWKS refresh for ~6-week rotation, Cloud Run ingress lock)		?	?	counsel (public signup)
P4a	**ONNX export of WASB** + implement the ABC <code>`ONNXRunner`</code> (identical CSV) + owner byte-parity/golden gate on a real GPU ? **dissolves the dual-env**		?	?	owner GPU
P4b	**Training-in-UI decision** (durable out-of-process long-job runner + progress/cancel, surface <code>`gate_verdict ship=False`</code>) ? **recommend DEFER from v1**		?	Owner	?
P4c	**Annotation decision** ? bundle VLC 3.0.x plugin vs build the **in-browser HTML5 annotator** emitting the same rally CSV ? **recommend in-browser** , owner-gated		?	Owner	?

8. Open questions ? owner / external (blocking gates first)

1. ****01 ? Re-ratify L5/D10 vs the launcher+installer:**** is a Start-Menu shortcut + friendly launcher (opens **your** browser) + Windows installer acceptable as "a self-hosted webserver you spin up", or does any OS-menu affordance cross the "NOT a desktop app" line? ****Need it in writing before P2.**** (Reject ``pystray`/LGPL` regardless; `stdlib`webbrowser`` or an MIT tray lib only.)
2. ****02 ? Confirm v1 = LIGHT**** (Gemini/motion/cpu-mock, no torch/WASB) and that ****CUJ-6/7** (native WASB) is deliberately owner-only** until C3 + reproducible-config clear. ***This is the load-bearing scope call.***
3. ****03 ? WASB-C3 counsel sign-off:**** may the checkpoint be redistributed / auto-fetched into a public build at all? Determines whether a public HEAVY tier is **ever** possible vs WASB-on-cloud-only.
4. ****04 ? Cloud in a downloaded app = hosted-Khelsutra (none built) or BYO-GCP (the only thing built)?**** Decide before marketing the "push to cloud" toggle (D13 currently overstates hosted).
5. ****05 ? Training in v1 UI?**** Recommend defer (synth=mock, real=GPU+dual-env, ``ship=False``).
6. ****06 ? Annotation in v1?**** in-browser HTML5 annotator (recommended) vs bundle VLC vs defer.
7. ****07 ? Ship AI-drafted kn/hi to real users before native-speaker review, or gate the trilingual UI on review?****
8. ****08 ? Shipped DEFAULT `config.json` / profile**** = truly-local (so the box never surprise-bills on Gemini nor returns 0 rallies on WASB-without-key)?
9. ****09 ? Code-signing budget**** (Windows Authenticode now, macOS Developer-ID later) so non-tech users don't hit SmartScreen/AV prompts.

9. Definition of done

- [] ****01/02/08 decided in writing**** (scope + delivery shape + default profile).
- [] P0a?P0e merged: launching ``indexer --server`` from ****any**** CWD finds the same config/DB; ``/api/config`` leaks nothing; all CI lanes (incl. the new license-scan + built-artifact lanes) green at the coverage floor.
- [] P1a/P1b merged: FastAPI serves the ****real SPA**** single-origin with refresh-survival; capabilities/videos/reels return real data; contract test guards SPA?engine skew.
- [] P2a?P2d merged: a ****clean Windows VM ? installer ? reel downloaded (en+hi)**** with ****no Python/git/ffmpeg pre-installed****; uninstall is clean; upgrade preserves data.
- [] ****Every completed CUJ in the map below verified preserved**** (CUJ-6/7 explicitly owner-only and flagged, not silently dropped).
- [] HEAVY/cloud (P3) + ONNX (P4) tracked but ****not gating v1****.

CUJ preservation map ``[verified inventory]``

CUJ	Status	How preserved in the app	
-----	-----	-----	
CUJ-1	Local Gemini ? rallies ? reel (en+hi)	DONE LIGHT ships Gemini + full	
<code>`/api/process?jobs?query?compile?highlights`</code> on the in-process worker, storage=local; verbatim, needs key + resolved ffmpeg.			
CUJ-2	In-UI <code>`gs://`</code> bucket ingest	DONE <code>`acquire_local_input`</code> + gcs backend already accept <code>`gs://`</code> ; preserved when the user opts into a cloud storage backend (+ ADC).	
CUJ-3	Per-rally selection + query bar	DONE Pure SPA + <code>`/api/query`</code> + <code>`/api/compile`</code> ; preserved by bundling the built SPA single-origin + the contract test.	
CUJ-4	Reel download-to-local	DONE <code>`/api/compile`</code> + <code>`GET /api/highlights/...`</code> unchanged; resolved ffmpeg provides the encoder; app-home output dir keeps reels findable.	
CUJ-5	Trilingual SPA (en+kn+hi), offline	DONE Catalogs + OFL Noto fonts compiled into the bundle (no CDN); air-gap renderable. <i>*(kn/hi native review pending ? 07.)*</i>	
CUJ-6/7	Worker GCS?GPU WASB?GCS + cold-start dual-env	DONE (engine) **HEAVY/cloud tiers ONLY** (owner + provisioned GPU) via micromamba WASB env / Cloud Run image ? **deliberately absent from the LIGHT default** ; gated on C3 + reproducible config.	
CUJ-8	CLI single-shot (<code>`indexer --video-path --segmenter`</code>)	DONE Console + module endpoints untouched by packaging; available every tier (segmenter set limited by tier).	
CUJ-9	Annotate ? train	PARTIAL <code>`worker train`</code> CLI + rally-annotator CSV bridge remain runnable; honesty fences (synth=mock, <code>`ship=False`</code>) must surface in any UI; recommend **deferred** from v1 UI (P4b/c).	

10. References

****Internal code anchors ``[verified]``:** ``pyproject.toml:80-82,8-37,63-70``;
``backend/main.py:10,64-74,144-146,275-315,646,729-736,756``;
``backend/config/loader.py:30,112,202``; ``backend/config/presets.py``;
``backend/compute/__init__.py``; ``backend/storage/*``;

```
`backend/pipeline/detectors/native_wasb_runner.py:17,89,110-112`;
`backend/pipeline/stitcher/compiler.py:253`; `backend/features/rally_seq.py:17`;
`backend/infrastructure/database.py:88-110` (user_version=6); `training/gen0/harness.py:174-177`
(`ship=False`); `deploy/cloudrun/Dockerfile:33-37,54`. Internal docs:
`docs/COMPUTE_DECOUPLED_SERVING/`, `docs/DESKTOP_APP_PLAN.md` (superseded),
`docs/UI_PROPOSAL.md`, `docs/I18N_PLAN.md`, `kheIsutra/docs/{LOCAL_APPLIANCE_ARCHITECTURE,ANNO
TE_AND_TRAIN,DRY_RUN_CUJ,FRIENDLY_AND_MULTILINGUAL}.md`. External `[researched]`: uv torch-
index docs; Frigate/Ollama/ComfyUI/LM Studio distribution; conda-forge/micromamba ToS; ONNX
Runtime; ffmpeg GPL-vs-LGPL.
```

Changelog

- 2026-06-21 ? Created from an exhaustive research workflow (29 agents) + per-claim re-verification at file:line. Status DRAFT; scope call 02 + delivery shape 01 are the load-bearing owner gates.

36. user (2026-06-21T14:16:18.232Z)

File created successfully at: C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md (file state is current in your context ? no need to Read it back)

37. assistant (2026-06-21T14:16:55.265Z)

Done. I synced, ran an exhaustive 29-agent research workflow, **re-verified every load-bearing claim at `file:line` myself**, and wrote it up as a tracked project doc:
docs/PACKAGING_PLAN.md (untracked/uncommitted ? I won't open a PR unless you ask). Here's the plan.

The core idea

One pip-installable artifact whose behaviour is selected entirely by the decoupled-compute seam you already merged (`deployment.profile` / `compute\_target` / Storage+Compute backends), fronted by a friendly launcher + a Windows installer. **Not** a frozen monolith, **not** an Electron/Tauri shell. Local-vs-cloud is genuinely **one config flip** ? you built that seam (M1?M8); this is the **distribution layer** on top, not a rewrite.

The constraint that forces the shape

Two facts are load-bearing and verified:

- torch** is deliberately out of `pyproject` (CUDA/CPU wheels need `--index-url`, which PEP 621 can't express) ? `pip install .` gives a torch-less install.
- Native WASB** needs a **second, incompatible torch env** (conda py3.8 + torch 1.11+cu113) shelled out via `\$WASB\_PYTHON` (`native\_wasb\_runner.py:89`), while the main engine is torch 2.6.

? **No single venv or freeze can deliver the GPU-quality path.** So the honest v1 is a **LIGHT tier** (Gemini + motion + cpu-mock, **zero torch/GPU**) that preserves CUJ-1..5/8 and runs on any CPU box. Native-WASB (CUJ-6/7) ships **owner-only** until two blockers clear.

Three tiers

Tier	Who	Contents	Install
------	-----	----------	---------

---	---	---	---
-----	-----	-----	-----

LIGHT (v1 default)	Non-tech coach on Windows; any CPU	wheel + bundled SPA + Gemini/motion/cpu-mock; no torch/weights; ffmpeg via winget-offer	Inno Setup .exe ? launcher ? auto-opens `localhost:8000`; uv script for power users
---------------------------	------------------------------------	---	--

HEAVY/local-GPU	Owner + opt-in power user	+ micromamba WASB env via `\$WASB_PYTHON` + SHA-256-pinned weights	opt-in "Enable GPU" step; falls back to LIGHT
------------------------	---------------------------	--	---

CLOUD/server	GCP / BYO-GPU-VM	the `deploy/cloudrun` image (micromamba-ified) behind Cloudflare Access	`gcloud run jobs create` / VM scripts
---------------------	------------------	---	---------------------------------------

How the three hard problems get solved

- **Dual-torch:** v1 carries **no torch**. HEAVY uses a separate **micromamba+conda-forge** env (BSD ? also fixes the **Anaconda-ToS** trap in `Dockerfile:33-37`). **Strategic fix:** export WASB to ONNX ? **onnxruntime** (MIT, tens of MB, CPU+CUDA+DirectML) **dissolves the dual-env entirely** and the **DetectorRunner** ABC already names an **ONNXRunner**.

- ****Weights:**** never bake them in; fetch on first run, SHA-256-pinned + resumable. ****Public** download must NOT auto-fetch until counsel clears WASB-C3** (`gate\_verdict` already hardcodes `ship=False`).
- ****ffmpeg:**** don't redistribute (the encoder is GPL `libx264`); one resolver helper + a `winget` install Gyan.FFMpeg` first-run offer.
- ****SPA-in-another-repo:**** build `kheisutra/web` at package time, vendor `web/dist/` into the wheel (hatchling force-include) served single-origin. The SPA already uses same-origin `BASE='/api'` ? ****zero SPA code change****, and it collapses the appliance reverse-proxy into one FastAPI process.

Phased plan (small, default-OFF, reversible PRs)

- ****P0 ? engine prerequisites (Claude-doable, no behaviour change):**** ? ****app-data HOME**** threaded through config/`.env`/DB/output/static (today all CWD-relative, and `doctor.py` writes `config.json` to a *different* place than the server reads ? a silent data-loss footgun); ? ****redact `/api/config`**** (returns the whole `AppConfig` unauthenticated today); ? config-driven bind/CORS; ? declare ****scipy**** (undeclared serving dep) + license-scan + built-artifact CI lanes; ? single-instance lock.
- ****P1 ? UI bundling + the 3 net-new endpoints**** (`/api/capabilities`, `/api/videos` list, `/api/reels`) the wizard needs (confirmed absent) + `/api/version` contract test.
- ****P2 ? LIGHT installer + first-run wizard**** (key/compute/videos, en/kn/hi) + DB backup-before-migrate + uninstall manifest.
- ****P3 ? HEAVY/cloud (largely owner-gated):**** micromamba, weight fetcher (C3-gated), "push to cloud" toggle (confirm + INR estimate), M9 edge-auth.
- ****P4 ? ONNX export (dissolves dual-env) + training/annotation decisions**** (recommend deferring training-in-UI; recommend an in-browser annotator over bundling VLC).

****Every completed CUJ is mapped**** in §9 of the doc ? CUJ-6/7 (native WASB) is the **one** deliberately absent from the LIGHT default, flagged explicitly rather than silently dropped (the honest version of "preserve every CUJ").

The decisions I need from you (load-bearing, before P2)

The biggest two ? I'd value a quick answer since they gate everything downstream:

38. user (2026-06-21T14:19:49.229Z)

Your questions have been answered: "Is v1 = the LIGHT tier (Gemini/motion/cpu-mock, no torch/GPU/WASB), with native-WASB (CUJ-6/7) shipping owner-only until WASB-C3 counsel + a reproducible config clear?"="1, but with a clear north star to include WASB and all as the sole completion goal", "Does a Windows installer + Start-Menu shortcut + a launcher that opens YOUR browser at localhost satisfy D10 ('self-hosted webserver, NOT a desktop app'), or does any OS-menu affordance cross that line?"="I would want the plan to scale to Windows, Mac and Linux as well so we can go with 1 and spinning up a webserver to power the UI in all three cases is PERFECTLY GOOD! No need for any native component.". You can now continue with these answers in mind.

39. assistant (2026-06-21T14:20:56.466Z)

Two big calls resolved ? and both reshape the plan. Let me fold them in:

1. ****v1 = LIGHT first, but with the full WASB path as the *sole completion goal* (north star)**** ? not deferred-indefinitely, it's the DONE target.
2. ****Cross-platform (Win/Mac/Linux), webserver-spin-up powers the UI on all three, zero native component.**** This drops the Windows-only Inno-Setup-as-primary framing.

Let me update the doc to lock these.

40. user (2026-06-21T14:20:59.027Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

41. user (2026-06-21T14:21:06.590Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been

updated successfully. (file state is current in your context ? no need to Read it back)

42. user (2026-06-21T14:21:21.366Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

43. user (2026-06-21T14:21:31.182Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

44. user (2026-06-21T14:21:36.698Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

45. user (2026-06-21T14:21:44.055Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

46. user (2026-06-21T14:21:52.901Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

47. user (2026-06-21T14:22:01.964Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

48. user (2026-06-21T14:22:09.390Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

49. assistant (2026-06-21T14:22:28.716Z)

Both decisions are now locked into the doc. Here's what changed:

- **L9 (new, locked):** Cross-platform **Windows + macOS + Linux**, **no native component** ? the same local webserver powers the UI on all three; distribution is `uv/pip` + a **pure-Python launcher** (`indexer --app` ? boot server ? open your browser). Electron/Tauri/tray and native-installer-as-runtime are all out. OSes differ **only** in how ffmpeg/torch are acquired (winget / brew / apt).
- **L10 (new, locked):** **The full quality path is the sole completion goal.** **LIGHT** is the first increment; the doc's Definition of Done now splits **"v1 shippable"** from **"north-star DONE"**, and the latter explicitly requires native WASB (in the **default** download), BYO training, and in-app annotation ? P3/P4 are now **completion-gating**, not optional.
- The old Inno-Setup-Windows-first framing (P2b, \$4.5 tiers) is replaced by the cross-platform webserver-spin-up path; code-signing drops to "only if we ever ship clickable binaries."

Let me record this as a memory checkpoint so a future session doesn't relitigate the locked calls or spiral on the (now-resolved) scope question.

50. assistant (2026-06-21T14:22:48.261Z)

name: packaging-app-plan

description: NEW owner-initiated track (2026-06-21) ? ship the engine as a downloadable batteries-included Python app (local web service for inference/training/annotation, local + cloud). Tracked doc + 01/02 locked.

```

metadata:
  node_type: memory
  type: project
---
```

****Owner asked (2026-06-21):**** bake the highlight indexer into a ****downloadable, batteries-included Python app**** ? a user downloads it (NOT `git clone`/build-from-source) and it ****spins up a local web service**** for inference + training + (optional) rally annotation, runnable ****locally AND at cloud services****. Preserve every completed CUJ.

****Artifact:**** `docs/PACKAGING\_PLAN.md` (engine repo) ? tracked project doc, Status DRAFT. Built from a 29-agent research workflow + per-claim re-verification at file:line. §7 progress tracker = source of truth (P0?P4).

****The shape (build ON the merged decoupled-compute seam, don't reinvent):**** ONE pip wheel, profile-selected (`deployment.profile`/`compute\_target`); the built ****khelsutra`web` SPA vendored into the wheel**** served single-origin by FastAPI (SPA already uses same-origin `BASE`/`api` ? zero SPA change). Three tiers: ****LIGHT**** (Gemini/motion/cpu-mock, NO torch) / ****HEAVY**** (native WASB via a micromamba `\$WASB\_PYTHON` env) / ****CLOUD**** (the `deploy/cloudrun` image behind Cloudflare Access).

****The hard constraint:**** torch is out-of-band (`--index-url`) AND native WASB needs a SECOND incompatible torch env (1.11) ? no single venv/freeze delivers the GPU path. Strategic fix = ****export WASB to ONNX**** (onnxruntime MIT) which dissolves the dual-env. ffmpeg is GPL `libx264` ? don't redistribute (resolver + per-OS install). WASB ****weights are C3-unresolved**** (`gate\_verdict ship=False`) ? public download must NOT auto-fetch until counsel clears.

****OWNER DECISIONS LOCKED (don't relitigate):****

- ****01/L9:**** cross-platform ****Win+Mac+Linux****, ****NO native component**** ? webserver-spin-up powers the UI on all 3; pure-Python launcher; reject Electron/Tauri/LGPL-tray/native-installer-as-runtime.
- ****02/L10:**** ****LIGHT-first**, but the FULL path (native WASB + BYO training + in-app annotation) is the SOLE completion goal / north star.****** DoD splits "v1 shippable" from "north-star DONE".

****Remaining gates:**** 03 WASB-C3 counsel (can weights ship in a public build at all?) · 04 hosted-khelsutra-cloud (none built) vs BYO-GCP · 05 training-in-v1 (recommend defer) · 06 annotation (recommend in-browser HTML5 annotator) · 07 kn/hi native review before public · 08 default profile=truly-local · 09 code-signing budget (only if clickable binaries).

****Claude-doable NOW (default-OFF, behaviour-neutral), if owner says go ? §7 P0:**** app-data HOME (today config/.env/DB/output are CWD-relative + `doctor.py` write disagrees with server read ? silent data-loss footgun) · redirect `GET /api/config` (returns whole AppConfig unauthenticated `main.py:144`) · config-driven bind/CORS · declare `scipy` (undeclared serving dep `features/rally\_seq.py:17`) + license-scan + built-artifact CI lanes · single-instance lock. Then P1 = bundle SPA + the 3 net-new endpoints (`/api/capabilities`, `/api/videos` list, `/api/reels` ? confirmed ABSENT).

****Relation to other tracks:**** parallel to the `[[next-session-ux-localization]]` UX/l10n focus and the decoupled-serving M9+ batch; this is the DISTRIBUTION layer under khelsutra `LOCAL\_APPLIANCE\_ARCHITECTURE.md`; supersedes `DESKTOP\_APP\_PLAN.md` narrative. `[[monetization-plan]]` `[[compute-stack-gcp-only]]` `[[lesson-verify-engine-surface-before-scoping]]`

51. user (2026-06-21T14:22:48.293Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md (file state is current in your context ? no need to Read it back)

52. user (2026-06-21T14:22:56.302Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

53. user (2026-06-21T14:23:01.280Z)

<system-reminder>Warning: the file exists but is shorter than the provided offset (60). The file has 49 lines.</system-reminder>

54. user (2026-06-21T14:23:05.029Z)

```
38 - [Eval dual-set regression](eval-dual-set-regression.md) ? owner directive (2026-06-16,
proven practice): keep a **RAW eval set** (currently-bad videos, random sample) + a **GROWING
GOLDEN-finalized eval set** (promote raw?golden when at-par) + a separate **golden TRAINING
set**; every quality launch must NOT regress on EITHER eval set (don't lose a graduated win).
The **ablation layer A/B**s these. Checkpointed video folder; wire dual-set aggregate into
`run_regression.py` when R2 lands.
39
40 ## Monetization (hosted API)
41 - [Monetization plan](monetization-plan.md) ? hosted video?highlights API on a GPU VM;
licensing & per-video cost are first-class; user is Bangalore-based
42 - [Monetization research status](monetization-research-status.md) ? audit done:
Ultralytics AGPL swapped out (PR #8); MonoTrack/YOLO-NAS weights avoided; codec=watch; full
report in docs/
43 - [Machine-pooling idea](machine-pooling-idea.md) ? owner "food for thought"
(2026-06-20): pool GPU machines (queue-on-one-host vs Cloud-Run scale-to-zero) instead of spin-
up/kill per task ? born from a single-GPU-slot collision between two sessions (quota=1). Infra-
track design item.
44 - [UI-driven cloud-serving gap](ui-driven-cloud-serving-gap.md) ? **owner directive
(2026-06-21): the cloud highlight flow must be UI-driven + ALPHA-SHAREABLE, not SSH/CLI; surface
compute options (local vs cloud) in the UI at start.** Today = SSH-only; M12 ? the fix. Needs:
M7 deploy (persistent service) + SPA ingest screen (khelsutra B-P2) + M9 auth ON + a UI compute-
picker (D13/D15). Engine-makes-possible vs SPA-makes-usable.
45
46 ## Forward-looking
47 - [Sibling internship repo](sibling-internship-repo.md) ? khelacharya shot-intelligence
layer builds ON TOP of our rally segmentation; kept off master at branch level (NOT a separate
GitHub repo ? corrected 2026-06-06)
48 - [Paper co-author aim](paper-coauthor-aim.md) ? **standing side-aim**: co-author a
paper on "local, cheap & efficient rally detection via multiple clever ideas". Keep
`docs/QUALITY_ITERATIONS.md` as the ablation-ready empirical backbone; honest measured results
ONLY.
49
```

55. user (2026-06-21T14:23:12.216Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

56. assistant (2026-06-21T14:23:33.569Z)

The plan is complete and recorded. Summary of what I delivered:

****Artifacts****

- ****docs/PACKAGING_PLAN.md**** ? full tracked project doc (template-compliant: decisions-locked table, threat model, \$7 progress tracker P0?P4, CUJ-preservation map, DoD split into v1-shippable vs north-star-DONE). Untracked/uncommitted ? no PR opened.
- Memory checkpoint written so the two locked decisions don't get relitigated, cross-linked to the existing [ui-driven-cloud-serving-gap](ui-driven-cloud-serving-gap.md) directive.

****The plan in one breath:**** one pip wheel, profile-selected (reusing your merged decoupled-compute seam), with the khelsutra SPA vendored in and served single-origin; cross-platform Win/Mac/Linux via `uv`/pip + a pure-Python launcher, no native component; ****LIGHT** (Gemini/CPU, no torch) ships first** and preserves CUJ-1..5/8, while the ****full native-WASB + training + annotation path is the locked completion goal**** ? reached by micromamba in the short term and an ****ONNX export that dissolves the dual-torch problem**** long term. Weights stay owner-only until WASB-C3 counsel clears; ffmpeg is resolved per-OS, never redistributed (GPL `libx264`).

****Method note:**** the recommendations are workflow-derived but I re-verified every load-bearing

claim at `file:line` myself (the `/api/config` leak, the CWD-relative state footgun, the 3 missing endpoints, undeclared scipy, the GPL encoder, the Anaconda-ToS image, `ship=False`) ? none are taken on faith.

****Suggested first move (your call ? I haven't started, per the no-implementation-without-approval rule):**** the ****P0a app-data HOME**** PR. It's behaviour-neutral (default-OFF when launched from the repo root), it's the prerequisite for the installer/upgrade/uninstall story, and it independently ****fixes a real silent data-loss bug today**** ? `scripts/doctor.py` writes `config.json` to a different place than the server reads it from.

Want me to:

1. ****Start P0a**** (and the rest of P0) as small default-OFF PRs, or
2. take the ****remaining gates to you first**** (03 WASB-C3 counsel, 04 hosted-vs-BYO-cloud) since they shape the north-star path, or
3. leave it as a plan for now?

57. user (2026-06-21T14:26:01.317Z)

Start P0a (and the rest of P0) as small default-OFF PRs please. And as for hosting test, you are free to spin up GCP GPU VMs as long as the budget cap I have specified is respected. As mentioned, prefer a "proof and validation (via screencast/screenshots etc.) driven development"

58. assistant (2026-06-21T14:28:59.502Z)

On it. The owner's standing constraints make the **how** matter as much as the **what**: this is a shared checkout with a sibling session's uncommitted M9-auth WIP (`models.py`, `auth.py`) sitting in the tree, so I'll work in an ****isolated git worktree off `origin/master`**** (never disturbing that branch/WIP), stage explicit paths only, and run an adversarial review per PR. P0 needs no GPU, so I'll park your VM offer for the P3/P4 WASB validation (where I'll confirm cmd + cost and delete-after, per the cap).

Let me set up the isolated worktree and verify I can operate in it.

59. user (2026-06-21T14:29:05.126Z)

```
=== origin/master HEAD ===
7835f34 feat(gcp): always-install Ops Agent on GPU VMs + a cloud-GPU dry-run guide (#291)
=== existing worktrees ===
C:/Users/avidu/Projects/badminton-highlight-indexer
a1bec6f [fix/dryrun-guide-font-and-ui-tracking]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/agent-a771ff1d85e18f2ed
7e0881a [perf/wasb-fast-frames]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/bhi-execpolicy
5cb7bbf [bhi-execpolicy]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/boxhill17
8736e2d [run/boxhill17-highlights]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/gap-doc
f04f3f3 [docs/rally-gap-closure]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/hardchop
614325d [feat/w1-hard-cap-fallback]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/pr-stitch
a6adcc5 [fix/runner-stitch-watermark]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/recdoc
9fd4215 [docs/recording-single-match]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/resweep
fe34013 [docs/lovo-resweep-n13]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing
a3b209a [master]
C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/w2-safeguard
5ba6955 [feat/w2-never-empty-safeguard]
=== create isolated worktree off origin/master ===
Preparing worktree (new branch 'feat/packaging-p0a-app-home')
branch 'feat/packaging-p0a-app-home' set up to track 'origin/master'.
HEAD is now at 7835f34 feat(gcp): always-install Ops Agent on GPU VMs + a cloud-GPU dry-run
```

```
guide (#291)
=== verify worktree branch + clean tree ===
feat/packaging-p0a-app-home
7835f34 feat(gcp): always-install Ops Agent on GPU VMs + a cloud-GPU dry-run guide (#291)
```

60. assistant (2026-06-21T14:29:19.990Z)

origin/master has already moved to `7835f34` (#291) since my session start, and the main checkout flipped to a *different* branch (`fix/dryrun-guide-font-and-ui-tracking`) ? exactly the shared-checkout churn the protocol guards against. My worktree is correctly based on the fresh `origin/master`. Let me confirm file access in the worktree and read the clean code I'll modify.

61. user (2026-06-21T14:29:20.307Z)

```
1      """Configuration loader and saver.
2
3      Provides a single place to load validated AppConfig with built-in defaults
4      merged with optional file overrides (config.json).
5
6      This replaces scattered json.load + .get() patterns throughout the codebase.
7      """
8
9      from __future__ import annotations
10
11     import json
12     import logging
13     import os
14     from pathlib import Path
15     from typing import Any
16
17     from pydantic import BaseModel
18
19     from .models import AppConfig
20     from .presets import (
21         COMPUTE_TARGET_FOR_PROFILE,
22         PROFILE_PRESETS,
23         is_known_profile,
24         preset_for,
25         suggest_profile,
26     )
27
28     logger = logging.getLogger(__name__)
29
30     DEFAULT_CONFIG_PATH = Path("config.json")
31
32
33     def _deep_merge(base: dict[str, Any], override: dict[str, Any]) -> dict[str, Any]:
34         """Recursively overlay ``override`` onto ``base`` (override wins; dicts merge,
35         scalars
36         replace). Used ONLY for the profile-preset precedence path so a config.json sub-key
37         overrides a preset sub-key while sibling preset/default sub-keys survive ? the no-
38         profile
39         path keeps the original shallow ``{**defaults, **file_data}`` merge unchanged."""
40         out = dict(base)
41         for key, val in override.items():
42             if isinstance(val, dict) and isinstance(out.get(key), dict):
43                 out[key] = _deep_merge(out[key], val)
44             else:
45                 out[key] = val
46         return out
47
48     def _resolve_explicit_profile(file_data: dict[str, Any]) -> str:
49         """The EXPLICITLY chosen deployment profile, env winning over config.json. ``""`` if
50         none
```

```

49         is set. Auto-detect is deliberately NOT consulted here ? it only suggests (D15), so
a bare
50         environment never silently selects a profile (and never silently spends on cloud).
51
52         An *unrecognised* ``DEPLOYMENT_PROFILE`` env value warns and FALLS THROUGH to the
file's
53         profile (rather than suppressing it) ? so an env typo can't leave the recorded
profile
54         asserting a preset that was never applied. A bad value in config.json is left to
surface as
55         a hard, typed-Literal error downstream (config-file correctness)."""
56         env_profile = os.environ.get("DEPLOYMENT_PROFILE")
57         if env_profile and env_profile.strip():
58             ep = env_profile.strip()
59             if is_known_profile(ep):
60                 return ep
61             logger.warning(
62                 "[CONFIG] unrecognized DEPLOYMENT_PROFILE env value %r (valid: %s); ignoring
it "
63                 "and falling back to config.json deployment.profile.", ep, ",
".join(PROFILE_PRESETS),
64             )
65             # fall through to the file's profile (env typo must not suppress a valid file
choice)
66         dep = file_data.get("deployment")
67         if isinstance(dep, dict):
68             prof = dep.get("profile")
69             if isinstance(prof, str) and prof.strip():
70                 return prof.strip()
71         return ""
72
73
74     def _unknown_key_paths(data: dict[str, Any], model_cls: type[BaseModel], prefix: str =
"") -> list[str]:
75         """Dotted paths in `data` that the typed config will not honor.
76
77         Two silent failure modes exist without this check: a typo'd TOP-LEVEL key is
78         kept by AppConfig's ``extra="allow"`` but never read by typed code, and a
79         typo'd key INSIDE a section is silently dropped by Pydantic. Either way the
80         user's intent is lost ? collect every such key so load_config can warn.
81         """
82         unknown: list[str] = []
83         fields = model_cls.model_fields
84         for key, val in data.items():
85             if key not in fields:
86                 unknown.append(prefix + key)
87                 continue
88             ann = fields[key].annotation
89             if isinstance(val, dict) and isinstance(ann, type) and issubclass(ann,
BaseModel):
90                 unknown.extend(_unknown_key_paths(val, ann, f"{prefix}{key}.")
91         return unknown
92
93
94     def _get_builtin_defaults() -> dict[str, Any]:
95         """Return a plain dict of the current built-in defaults.
96
97         The AppConfig model is the single source of truth for defaults.
98         """
99         return AppConfig().model_dump(mode="json")
100
101
102     def load_config(path: str | Path | None = None) -> AppConfig:
103         """Load configuration with the following precedence:
104

```

```

105     1. Built-in defaults defined in the Pydantic models.
106     2. Values from the JSON file (if present and readable).
107     3. (Future) environment / CLI overrides.
108
109     Missing file or unreadable file ? returns a fully defaulted AppConfig.
110     Unknown keys in the file are tolerated (extra="allow" on the model).
111     """
112     cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
113
114     file_data: dict[str, Any] = {}
115     if cfg_path.exists():
116         try:
117             with open(cfg_path, "r", encoding="utf-8") as f:
118                 file_data = json.load(f)
119         except Exception as exc:
120             # Non-fatal: continue with defaults + whatever we could parse.
121             # In production we might log here.
122             logger.warning(f"Failed to read {cfg_path}: {exc}. Using defaults + partial
data.")
123
124     # A config.json that is valid JSON but not an OBJECT (e.g. `[ ]`, `42`, `"x"`,
`null`) would
125     # otherwise crash startup: a truthy non-mapping hits `.items()` in
_unknown_key_paths, and any
126     # non-mapping breaks the `{**defaults, **file_data}` unpack. Degrade to defaults +
warn, per the
127     # loader's "bad input is non-fatal" contract.
128     if not isinstance(file_data, dict):
129         logger.warning("%s is not a JSON object (got %s); ignoring it and using
defaults.",
130                        cfg_path, type(file_data).__name__)
131     file_data = {}
132
133     if file_data:
134         unknown = _unknown_key_paths(file_data, AppConfig)
135         if unknown:
136             logger.warning(
137                 "[CONFIG WARNING] %s contains keys the typed config does not "
138                 "recognize (typo'd or removed knobs ? top-level extras are kept "
139                 "but unread; unknown section keys are silently dropped): %s",
140                 cfg_path, ", ".join(sorted(unknown)),
141             )
142
143     # Precedence: built-in defaults < profile preset < config.json (< env/--set, applied
at
144     # consumption sites downstream). A *deployment profile* is opt-in ? it is applied
ONLY when
145     # explicitly selected (DEPLOYMENT_PROFILE env or config.json deployment.profile).
With no
146     # profile the merge is byte-identical to the pre-M1 loader, so the default path is
unchanged.
147     defaults = _get_built_in_defaults()
148     profile = _resolve_explicit_profile(file_data)
149
150     if profile and not is_known_profile(profile):
151         # Only an unrecognised value from config.json reaches here (env typos already
fell back
152         # in _resolve_explicit_profile). Warn + apply no preset; the bad value also hits
the
153         # typed Literal at model_validate below, a hard clear error ? loud, never
silent.
154         logger.warning(
155             "[CONFIG] unrecognized deployment profile %r (valid: %s); applying no
preset.",
156             profile, ", ".join(PROFILE_PRESETS),

```

```

157         )
158         profile = ""
159
160     if profile:
161         merged = _deep_merge(_deep_merge(defaults, preset_for(profile)), file_data)
162         # Record the profile actually applied ? the env may have chosen it over
163         config.json,
164         # so re-stamp it onto the merged deployment section to keep the record honest.
165         merged.setdefault("deployment", {})
166         if isinstance(merged["deployment"], dict):
167             merged["deployment"]["profile"] = profile
168             ct = merged["deployment"].get("compute_target")
169             canonical = COMPUTE_TARGET_FOR_PROFILE.get(profile)
170             if not ct and canonical:
171                 # An active profile with no explicit compute_target (e.g. config.json
172                 # to "") gets its canonical target back ? the record never silently lies
173                 # active profile's compute_target (it stays meaningful for the M2+
174                 ComputeBackend).
175                 merged["deployment"]["compute_target"] = canonical
176                 elif ct and canonical and ct != canonical:
177                     logger.warning(
178                         "[CONFIG] deployment profile %r normally uses compute_target=%r but
179                         "
180                         "config.json overrides it to %r ? honoring the explicit override.",
181                         profile, canonical, ct,
182                     )
183                     logger.info("[CONFIG] deployment profile ACTIVE: %s (compute_target=%s)",
184                                 profile, merged["deployment"].get("compute_target"))
185         else:
186             # No profile ? the original shallow merge (Pydantic re-defaults nested
187             sections).
188             merged = {**defaults, **file_data}
189             suggestion = suggest_profile() # cheap, env-only (no torch); SUGGESTS only
190             (D15)
191             if suggestion:
192                 logger.debug(
193                     "[CONFIG] no deployment.profile set (manual knobs). Environment suggests
194                     %r ? "
195                     "set deployment.profile or DEPLOYMENT_PROFILE to apply (auto-detect
196                     never "
197                     "auto-applies ? D15).", suggestion,
198                 )
199
200     return AppConfig.model_validate(merged)
201
202
203 def save_default_config(path: str | Path | None = None) -> Path:
204     """Write a fresh default config.json.
205
206     Used by scripts/doctor.py diagnostics and the --setup CLI path.
207     Overwrites if the file exists.
208     """
209     cfg_path = Path(path) if path is not None else DEFAULT_CONFIG_PATH
210     cfg = AppConfig()
211     cfg_path.parent.mkdir(parents=True, exist_ok=True)
212     cfg_path.write_text(cfg.model_dump_json(indent=4), encoding="utf-8")
213     return cfg_path
214
215
216 # Backwards-compatible alias used by some older call sites during transition.
217 def get_default_config() -> dict[str, Any]:
218     """Return defaults as a plain dict (for transitional use only)."""
219     return _get_builtin_defaults()

```

62. assistant (2026-06-21T14:29:24.220Z)

File access in the worktree works. Let me read the rest of the surfaces P0a threads through.

63. user (2026-06-21T14:29:25.137Z)

```

1      """Path-safety helpers for the API boundary (untrusted request fields)."""
2      import os
3      from typing import Any, Optional
4
5
6      def output_base(config: Any) -> str:
7          """The configured base dir for a segmenter's LOCAL scratch/staging artifacts
8              (COMPUTE_DECOUPLED_SERVING M3). Returns ``output.output_dir`` (default ``"output"``)
9              so the chunk-staging / AI-handoff / temp dirs follow ``--output-dir`` /
10             ``output.output_dir`` instead of a hardcoded ``./output``. The DEFAULT is ``"output"`` ? byte-identical to the prior
11             literal.
12             This is the LOCAL output dir, deliberately NOT the storage ``workspace_root`` (which
13             may be a ``gs://`` URI you can't ``mkdirs``). Accepts the typed AppConfig or the legacy
14             dict-shim
15             (``.get("output", {})).get("output_dir")``); any miss or empty value falls back to
16             ``"output"``. """
17             try:
18                 out = config.get("output", {})).get("output_dir")
19             except Exception:
20                 out = None
21             return out or "output"
22
23
24      def safe_output_filename(name: str) -> str:
25          """Validate an output filename is a BARE name (no dir part, no traversal, not
26             absolute).
27             ``os.path.join(output_dir, name)`` silently discards ``output_dir`` when ``name`` is
28             absolute,
29             and ``..`` segments escape it ? so an unchecked ``output_filename`` is an arbitrary-
30             write
31             primitive. Returns the name unchanged when safe; raises ValueError otherwise.
32             """
33             if not name or not name.strip():
34                 raise ValueError("output_filename must not be empty")
35             if "\x00" in name:
36                 raise ValueError("output_filename must not contain a null byte")
37             if "/" in name or "\\" in name:
38                 raise ValueError(f"output_filename must not contain a path separator: {name!r}")
39             if os.path.isabs(name) or name != os.path.basename(name) or name in (".", ".."):
40                 raise ValueError(f"output_filename must be a bare filename, got: {name!r}")
41             return name
42
43
44      def resolve_under_root(path: str, root: str) -> str:
45          """Return the real path, asserting it is contained under ``root`` (raises ValueError
46             if not). """
47             rp = os.path.realpath(path)
48             rr = os.path.realpath(root)
49             try:
50                 contained = os.path.commonpath([rp, rr]) == rr
51             except ValueError: # different drives on Windows
52                 contained = False

```

```

48         if not contained:
49             raise ValueError(f"path {path!r} escapes the allowed root {root!r}")
50         return rp
51
52
53     def validate_input_path(path: str, allowed_root: Optional[str] = None) -> None:
54         """Reject null bytes and (when an allowed_root is configured) out-of-root inputs.
55
56         With no allowed_root (default, single-user local mode) only the null-byte check
57         applies,
58         so the local 'point it at any of my videos' workflow is preserved. Setting an
59         allowed_root (a future hosted/tenant deployment) enforces containment.
60         """
61         if "\\x00" in path:
62             raise ValueError("video_path must not contain a null byte")
63         if allowed_root:
64             resolve_under_root(path, allowed_root)

```

64. user (2026-06-21T14:29:25.512Z)

```

1     import sys
2     import os
3     import subprocess
4     import shutil
5
6     # This diagnostics/bootstrap script lives in scripts/ but operates on the repo
7     # root (requirements.txt, config.json). Resolve the root from this file's
8     # location so it works regardless of the caller's cwd.
9     REPO_ROOT = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
10
11     def check_command(cmd):
12         """Check if a system command is available."""
13         return shutil.which(cmd) is not None
14
15     def run_pip_install():
16         """Install required Python packages."""
17         print("Installing Python packages from requirements.txt...")
18
19         # Check if NVIDIA GPU is present
20         has_gpu = check_command("nvidia-smi")
21
22         requirements_path = os.path.join(REPO_ROOT, "requirements.txt")
23         pip_cmd = [sys.executable, "-m", "pip", "install", "--user", "-r",
24         requirements_path]
25
26         import platform
27         if has_gpu:
28             print("[INFO] NVIDIA GPU detected. Configuring PyTorch to download CUDA-enabled
29             wheels.")
30             pip_cmd.extend(["--extra-index-url", "https://download.pytorch.org/whl/cu124"])
31         elif platform.system() == "Darwin":
32             # macOS has no CUDA; the default PyPI wheel already ships CPU + Apple MPS.
33             print("[INFO] macOS detected. Using default PyTorch wheels (CPU + Apple MPS).")
34         else:
35             print("[INFO] No NVIDIA GPU detected. Defaulting to PyTorch CPU wheels.")
36             pip_cmd.extend(["--extra-index-url", "https://download.pytorch.org/whl/cpu"])
37
38         try:
39             subprocess.check_call(pip_cmd)
40             print("[SUCCESS] Python packages installed.")
41             return True
42         except subprocess.CalledProcessError as e:
43             print(f"[ERROR] Failed to install Python packages: {e}")
44             return False

```

```

43
44 def init_default_config():
45     """Initialize a default config.json if it doesn't exist.
46
47     Now delegates to the canonical backend.config loader so that defaults
48     stay in sync with the Pydantic models (single source of truth).
49     """
50     try:
51         from backend.config import save_default_config
52     except Exception as e:
53         print(f"[ERROR] Could not import backend.config loader: {e}")
54         print("[INFO] Falling back to no-op for config initialization.")
55         return
56
57     config_path = os.path.join(REPO_ROOT, "config.json")
58     if os.path.exists(config_path):
59         print(f"[INFO] {config_path} already exists. Skipping initialization.")
60         return
61
62     try:
63         saved_path = save_default_config(config_path)
64         print(f"[SUCCESS] config.json initialized with default settings at {saved_path}.")
65     except Exception as e:
66         print(f"[ERROR] Failed to create config.json via loader: {e}")
67
68 def run_diagnostics():
69     """Perform system-wide checks for FFmpeg and PyTorch/CUDA."""
70     print("-" * 50)
71     print("SPORTS HIGHLIGHT INDEXER SETUP DIAGNOSTICS")
72     print("-" * 50)
73
74     # 1. Check Python Version
75     print(f"Python Version: {sys.version_info.major}.{sys.version_info.minor}.{sys.version_info.micro}")
76     if sys.version_info.major < 3 or (sys.version_info.major == 3 and sys.version_info.minor < 8):
77         print("[WARNING] Python 3.8+ is highly recommended!")
78     else:
79         print("[PASS] Python version check passed.")
80
81     print("-" * 50)
82
83     # 2. Check FFmpeg / FFprobe
84     ffmpeg_ok = check_command("ffmpeg")
85     ffprobe_ok = check_command("ffprobe")
86
87     if ffmpeg_ok:
88         print("[PASS] FFmpeg is installed and available in PATH.")
89     else:
90         print("[FAIL] FFmpeg is NOT found in your PATH.")
91
92     if ffprobe_ok:
93         print("[PASS] FFprobe is installed and available in PATH.")
94     else:
95         print("[FAIL] FFprobe is NOT found in your PATH.")
96
97     if not (ffmpeg_ok and ffprobe_ok):
98         import platform
99         print("\n" + "!" * 50)
100         print("CRITICAL: FFmpeg and FFprobe are REQUIRED to run this application.")
101         _sys = platform.system()
102         if _sys == "Windows":
103             print("Install on Windows: winget install Gyan.FFmpeg")
104         elif _sys == "Darwin":

```



```

105         print("Install on macOS:    brew install ffmpeg")
106     else:
107         print("Install on Linux:    sudo apt-get install -y ffmpeg    (or your
distro's package manager)")
108         print("Or download a build from https://ffmpeg.org/download.html and put
ffmpeg/ffprobe on PATH.")
109         print("!" * 50 + "\n")
110     else:
111         print("[SUCCESS] All essential media command line tools are ready.")
112
113     print("-" * 50)
114
115     # 3. Check ML / GPU (Optional but good)
116     try:
117         import torch
118         cuda_ok = torch.cuda.is_available()
119         if cuda_ok:
120             print(f"[PASS] PyTorch loaded successfully. CUDA (GPU) is AVAILABLE:
{torch.cuda.get_device_name(0)}")
121         else:
122             print("[INFO] PyTorch loaded successfully, but CUDA (GPU) is NOT available.
Falling back to CPU.")
123     except ImportError:
124         print("[INFO] PyTorch is not installed. Relevance checks will run using
lightweight visual heuristics.")
125
126     print("=" * 50)
127     return ffmpeg_ok and ffprobe_ok
128
129 if __name__ == "__main__":
130     init_default_config()
131     media_ok = run_diagnostics()
132     run_pip_install()
133     print("\n[INFO] Setup complete. If diagnostics passed, you can now run the app!")
134

```

65. user (2026-06-21T14:29:30.509Z)

```

1     import os
2     import sys
3     import argparse
4     import logging
5     import uuid
6     from urllib.parse import quote
7     import uvicorn
8     from dotenv import load_dotenv
9
10    load_dotenv()
11
12    # Centralized logging setup for the indexer (hot paths use logger, CLI output uses print
for direct user-facing messages).
13    logging.basicConfig(
14        level=logging.INFO,
15        format='%(asctime)s [%(levelname)s] %(name)s: %(message)s',
16        datefmt='%H:%M:%S'
17    )
18    logger = logging.getLogger(__name__)
19    from fastapi import FastAPI, HTTPException, Request
20    from fastapi.middleware.cors import CORSMiddleware
21    from pydantic import BaseModel
22    from typing import List, Dict, Any, Optional
23
24    from fastapi.staticfiles import StaticFiles
25    from fastapi.responses import FileResponse, HTMLResponse, JSONResponse
26    from backend.pipeline.validators.base import registry

```

```

27     from backend.infrastructure.database import Database
28     from backend.pipeline.segmenters.base import segmenter_registry
29     from backend.pipeline.stitcher.compiler import VideoCompiler
30     from backend.utils.hashing import compute_video_id # canonical file-identity hashing
31     from backend.utils.paths import resolve_under_root, safe_output_filename,
validate_input_path
32     from backend.config import ( # new typed config
33         load_config as load_app_config,
34         AppConfig,
35         StitchingConfig,
36         enforce_startup_checks,
37         StartupCheckError,
38     )
39     from backend.results import Failure # standardized error/result contract
40     from backend.reporting import emit_indexer_report
41     from backend import storage # M2: URI-aware input acquisition (local = identity
passthrough)
42     from backend import billing # M8: per-job cost ledger (USD internal / INR display)
43     from backend.api import auth as edge_auth # M9: Cloudflare Access edge-auth (default-
OFF)
44
45     app = FastAPI(title="Sports Highlight Indexer API")
46
47
48     def _cors_origins_from_env(env: Optional[Dict[str, str]] = None) -> List[str]:
49         """M9: CORS allow-origins from the ``RALLY_CORS_ALLOW_ORIGINS`` env var (comma-
separated),
50         default ``["*"]``. Read from the ENV (not a cwd ``config.json``) so importing
``backend.main``
51         does zero filesystem I/O ? a hosted/multi-tenant deploy sets this to lock CORS to
its origin."""
52         src = os.environ if env is None else env
53         raw = (src.get("RALLY_CORS_ALLOW_ORIGINS", "") or "").strip()
54         origins = [o.strip() for o in raw.split(",") if o.strip()]
55         return origins or ["*"]
56
57
58     # CORS for the web UI. allow_origins='' with allow_credentials=True is rejected by
browsers per the
59     # fetch spec + is unsafe to carry into multi-tenant; this tool uses no cookies, so
credentials stay
60     # off. The origin list is configurable via the env var above (default ["*"]) so a hosted
deploy can
61     # lock it to its origin; the middleware is fixed at startup.
62     app.add_middleware(
63         CORSMiddleware,
64         allow_origins=_cors_origins_from_env(),
65         allow_credentials=False,
66         allow_methods=["*"],
67         allow_headers=["*"],
68     )
69
70     # Serves static frontend files directly (now under api/ per approved structure)
71     os.makedirs("backend/api/static", exist_ok=True)
72     app.mount("/static", StaticFiles(directory="backend/api/static"), name="static")
73
74     @app.get("/", response_class=HTMLResponse)
75     def serve_index():
76         index_path = "backend/api/static/index.html"
77         if os.path.exists(index_path):
78             with open(index_path, "r", encoding="utf-8") as f:
79                 return f.read()
80         return "<h3>Sports Highlight Indexer Server is Running. Please create index.html in
backend/api/static/</h3>"
81

```

```

82     # Global variables initialized at startup
83     db_instance: Optional[Database] = None
84     global_config: Optional[AppConfig] = None # now typed (Pydantic model)
85
86     # Async job worker ? drain/wake loop (DEFERRED_HARDENING_PLAN #16, EXECUTION_POLICY.md
87     # now lives in backend.api.job_worker. Re-exported here so the names stay addressable on
88     # `backend.main` (`JobWaiting`, `_drain_due_jobs`, `_arm_wake_timer`, `_get_executor`,
89     # `_run_claimed_job`, `_MAX_DRAIN_WAKE_SEC`, ...). The worker resolves these seams
THROUGH
90     # this module object at call-time, so `monkeypatch.setattr(backend.main, ...)` still
91     # intercepts the inter-function calls (a parked job arms the PATCHED `_arm_wake_timer`,
92     # so no real daemon timer leaks in tests). Behavior is unchanged by the extraction.
93     from backend.api.job_worker import ( # noqa: F401 (re-exports: names stay addressable
on backend.main)
94         JobWaiting,
95         _arm_wake_timer,
96         _drain_due_jobs,
97         _get_executor,
98         _job_to_request,
99         _MAX_DRAIN_WAKE_SEC,
100         _maybe_record_job_cost,

```

66. assistant (2026-06-21T14:29:49.665Z)

The worktree shows ****M9 edge-auth has already merged to master**** (`main.py:43`, plus `\_cors\_origins\_from\_env` reading `RALLY\_CORS\_ALLOW\_ORIGINS`) ? so the `RALLY\_\*` env convention is established (my `RALLY\_APP\_HOME` fits) and P0c's CORS knob partly exists. Let me read the server-start path (config load, DB path, bind) and the `/api/config` handler at this fresh HEAD.

67. user (2026-06-21T14:29:50.868Z)

```

140         db_instance.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
141
142
143     # --- API Endpoints ---
144     @app.get("/api/config")
145     def get_config():
146         return global_config
147
148     @app.get("/api/health")
149     def health():
150         """Cheap liveness/readiness probe (no auth) for the SPA's offline banner
151         (HOSTING_PLAN.md:42). Reports the configured sport + default segmenter so the UI can
152         confirm the engine is reachable and which detector is wired. Never raises."""
153         sport = getattr(global_config, "sport", None)
154         indexing = getattr(global_config, "indexing", None)
155         return {
156             "status": "ok",
157             "sport": sport,
158             "default_segmenter": getattr(indexing, "default_segmenter", None),
159         }
160
161     def _execute_processing(req: ProcessRequest, progress=None) -> Dict[str, Any]:
162         """The full hash ? validate ? segment ? report pipeline for one video.
163
164         This is the former synchronous /api/process body, unchanged in behavior, now run on
165         the job worker thread (DEFERRED_HARDENING_PLAN #16). Returns the same response dict
166         the sync endpoint used to return (UI compatibility); the per-video observability
167         report is still emitted in `finally:` and grafted as response["reports"].
168
169         `progress` is an optional callable(stage: str) used to surface coarse job progress.

```

68. assistant (2026-06-21T14:29:55.499Z)

`/api/config` still leaks the full config (P0b stands). Now the server-start path ? config load, DB path, and bind.

69. user (2026-06-21T14:29:58.614Z)

```
=== db path / output_dir / load_config in main.py ===
33: load_config as load_app_config,
106:# New code should import load_app_config / AppConfig from backend.config directly.
107:def load_config(config_path: str = "config.json") -> Dict[str, Any]:
108:     cfg = load_app_config(config_path)
270:         # report.output_dir). Never raises. Surface the paths in the response.
394:     (which only writes into output_dir; the browser previously got back an unreachable
402:     output_dir = getattr(getattr(global_config, "output", None), "output_dir", "output") or
"output"
403:     path = os.path.join(output_dir, name)
405:         path = resolve_under_root(path, output_dir)
473:     output_dir = getattr(getattr(global_config, "output", None), "output_dir", "output") or
"output"
474:     compiler = VideoCompiler(output_dir, begin_padding=stitching.begin_padding,
498:     cfg = load_config() # legacy wrapper still works, returns dict for now
580:     parser.add_argument("--output-dir", default="output", help="Directory where processed
clips/highlights are saved")
621:         out_dir = args.output_dir
631:         stitching = load_app_config().stitching
650:     global_config = load_app_config() # returns typed AppConfig
653:     # The CLI --output-dir overrides the configured output directory. It now lives
654:     # on the typed model (OutputConfig.output_dir), so every consumer ? including
656:     global_config.output.output_dir = args.output_dir
657:     os.makedirs(global_config.output.output_dir, exist_ok=True)
659:     db_path = os.path.join(global_config.output.output_dir, global_config.output.db_name)
660:     db_instance = Database(db_path)
664:     swept = db_instance.fail_stale_jobs()
732:         print(f"[SUCCESS] Process complete. Indexed {len(segments)} play segments
inside database: {db_path}")
756:         uvicorn.run(app, host="127.0.0.1", port=args.port)

=== OutputConfig in models.py ===
230:class OutputConfig(BaseModel):
232:     db_name: str = "sports_indexer.db"
235:     output_dir: str = "output"
319:     # Output root for produced artifacts. "" => fall back to output.output_dir (non-
breaking).
337:     # Durable workspace root URI. "" => precedence: output_root, then output.output_dir.

=== Database constructor ===
12:     def __init__(self, db_path: str):
13:         self.db_path = db_path
34:         conn = sqlite3.connect(self.db_path, timeout=30.0)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

70. assistant (2026-06-21T14:30:02.562Z)

Let me read the full server-start block (575?760) to make surgical changes.

71. user (2026-06-21T14:30:03.597Z)

```
575
576
577     def main():
578         parser = argparse.ArgumentParser(description="Sports Highlight Indexer
Orchestrator")
579         parser.add_argument("--video-path", help="Local path to the MP4 sports video file")
580         parser.add_argument("--output-dir", default="output", help="Directory where
processed clips/highlights are saved")
```

```

581     parser.add_argument("--setup", action="store_true", help="Launch dependency checker
and setup utility")
582     parser.add_argument("--debug-clip", type=float, help="Timestamp (in seconds) of a
video clip to run detailed diagnostics on")
583     parser.add_argument("--server", action="store_true", help="Run the FastAPI local API
server")
584     parser.add_argument("--port", type=int, default=8000, help="Port to run the FastAPI
server on")
585     parser.add_argument("--max-frames", type=int, help="Limit processing to the first N
sub-sampled frames for debugging")
586
587     available_segmenters = list(segmenter_registry.list_available().keys())
588     parser.add_argument("--segmenter", choices=available_segmenters, help=f"Choose rally
detection segmenter. Available: {available_segmenters}")
589     parser.add_argument("--trim-start", type=float, help="Start time in seconds for the
debug trim segment feature")
590     parser.add_argument("--trim-end", type=float, help="End time in seconds for the
debug trim segment feature")
591     parser.add_argument("--trim-output", help="Output filename for the debug trimmed
segment (saves in output-dir unless absolute path)")
592
593     args = parser.parse_args()
594
595     if args.setup:
596         # Invoke the diagnostics/bootstrap script (renamed from the old root
597         # setup.py to scripts/doctor.py so it no longer shadows the PEP 517
598         # build system). Resolve its path from this file so `indexer --setup`
599         # works regardless of the caller's cwd.
600         import subprocess
601         doctor_path = os.path.join(
602             os.path.dirname(os.path.abspath(__file__)),
603             "scripts", "doctor.py",
604         )
605         print(f"[INFO] Invoking diagnostics: {doctor_path} ...")
606         subprocess.run([sys.executable, doctor_path])
607         return
608
609     if args.trim_start is not None and args.trim_end is not None:
610         if not args.video_path:
611             print("[ERROR] Please provide --video-path to extract a segment.")
612             return
613
614         # Determine output path
615         out_filename = args.trim_output or "trimmed_segment.mp4"
616         if os.path.isabs(out_filename):
617             out_path = out_filename
618             out_dir = os.path.dirname(out_path)
619             out_name = os.path.basename(out_path)
620         else:
621             out_dir = args.output_dir
622             out_path = os.path.join(out_dir, out_filename)
623             out_name = out_filename
624
625         os.makedirs(out_dir, exist_ok=True)
626         print(f"[CLI] Extracting debug segment from {args.trim_start}s to
{args.trim_end}s from {args.video_path}...")
627
628         # global_config isn't initialized this early; load the typed config so the trim
629         # clip carries the same configured stitching paddings + watermark as a real
630         # highlight (watermark default-on).
631         stitching = load_app_config().stitching
632         compiler = VideoCompiler(out_dir, begin_padding=stitching.begin_padding,
633                                 end_padding=stitching.end_padding,
634                                 watermark=stitching.watermark)
635         try:

```

```

635         res_path = compiler.compile_highlights(args.video_path, [(args.trim_start,
args.trim_end)], out_name)
636         print(f"[SUCCESS] Trimmed segment saved to: {res_path}")
637     except Exception as e:
638         print(f"[ERROR] Failed to compile trimmed segment: {e}")
639     return
640
641     if args.debug_clip is not None:
642         if not args.video_path:
643             print("[ERROR] Please provide --video-path to debug a specific clip.")
644             return
645         run_cli_diagnose_clip(args.video_path, args.debug_clip)
646     return
647
648     # Initialize Global Config and DB
649     global global_config, db_instance
650     global_config = load_app_config() # returns typed AppConfig
651     _enforce_startup_or_exit(global_config) # M5: fail fast on an incoherent ACTIVE
profile
652
653     # The CLI --output-dir overrides the configured output directory. It now lives
654     # on the typed model (OutputConfig.output_dir), so every consumer ? including
655     # /api/compile ? reads the same value instead of a write-only runtime hack.
656     global_config.output.output_dir = args.output_dir
657     os.makedirs(global_config.output.output_dir, exist_ok=True)
658
659     db_path = os.path.join(global_config.output.output_dir,
global_config.output.db_name)
660     db_instance = Database(db_path)
661
662     # Crash recovery (#16): jobs left queued/running by a previous process are dead ?
663     # the executor is in-process. Mark them failed so pollers see a terminal state.
664     swept = db_instance.fail_stale_jobs()
665     if swept:
666         logger.warning(f"[JOBS] Swept {swept} stale job(s) from a previous run ?
failed.")
667
668     # CLI processing mode (no web UI needed)
669     if args.video_path and not args.server:
670         print(f"[CLI] Processing video file: {args.video_path}")
671         # M2: bare/local path = identity (today's behaviour); a bucket/Drive URI is
fetched to scratch.
672         # The local-existence fast-fail only applies to a local input (stat the RESOLVED
key, not the
673         # raw URI); a remote fetch fails loudly. An unsupported scheme is a clean error,
not a crash.
674         storage_cfg = getattr(global_config, "storage", None)

```

72. assistant (2026-06-21T14:31:38.948Z)

I have the complete flow. Before writing, two checks: how heavy the package `\_\_init\_\_` is for the early `load\_dotenv` import, and every `DEFAULT\_CONFIG\_PATH` consumer repo-wide.

73. user (2026-06-21T14:31:45.193Z)

```

=== backend/config/__init__.py ===
"""Configuration package.

```

Public API:

```

from backend.config import load_config, AppConfig, Config

```

This package is the canonical place for all configuration access.
"""

```

from .loader import load_config, save_default_config

```

```

from .models import ( # re-export common ones
    AppConfig,
    Config,
    DeploymentConfig,
    IndexingConfig,
    StitchingConfig,
    StorageConfig,
    ValidationConfig,
)
from .presets import PROFILE_PRESETS, probe_cuda_available, suggest_profile
from .startup import (
    StartupCheck,
    StartupCheckError,
    enforce_startup_checks,
    run_startup_checks,
)

__all__ = [
    "load_config",
    "save_default_config",
    "AppConfig",
    "Config",
    "DeploymentConfig",
    "IndexingConfig",
    "StitchingConfig",
    "StorageConfig",
    "ValidationConfig",
    "PROFILE_PRESETS",
    "suggest_profile",
    "probe_cuda_available",
    "StartupCheck",

```

```

=== backend/utils/__init__.py (size) ===
1 backend/utils/__init__.py

```

Utils package

```

=== DEFAULT_CONFIG_PATH usage repo-wide ===
backend/config/loader.py:30:DEFAULT_CONFIG_PATH = Path("config.json")
backend/config/loader.py:112:     cfg_path = Path(path) if path is not None else
DEFAULT_CONFIG_PATH
backend/config/loader.py:202:     cfg_path = Path(path) if path is not None else
DEFAULT_CONFIG_PATH

```

```

=== app_home / RALLY_APP_HOME pre-existing? ===
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

74. assistant (2026-06-21T14:32:17.692Z)

`backend/utils/\_\_init\_\_.py` is empty (cheap early import), `DEFAULT\_CONFIG\_PATH` has **zero external consumers**, and no prior app-home code exists. Design locked: a stdlib-only `backend/utils/app\_home.py` (so the `load\_dotenv` import at `main.py:10` stays pydantic-free), defaulting to **CWD** when `RALLY\_APP\_HOME` is unset ? byte-identical to today². Writing the module + tests now.

75. assistant (2026-06-21T14:32:37.662Z)

"""Resolved application-data HOME for a packaged, batteries-included install.

PACKAGING_PLAN.md P0a. Today every piece of mutable state ? ``config.json``, ``.env``, the SQLite DB, ``output/`` ? is resolved RELATIVE TO THE CURRENT WORKING DIRECTORY, and ``scripts/doctor.py`` writes ``config.json`` to the repo root while the server reads it from the CWD. Launched from an OS app menu (CWD ? repo root) those disagree and the user silently "loses" their index. This module introduces ONE resolver every state path roots on.

Design contract ? ****DEFAULT IS BYTE-IDENTICAL TO TODAY****:

```
* ``RALLY_APP_HOME`` env set ? that directory is the app home (the packaged launcher sets it).
```

```
* ``RALLY_APP_HOME`` unset ? ``Path(".")`` i.e. the CWD, exactly the pre-P0a behaviour.
```

The OS-default location (``%APPDATA%`` / ``~/Library/Application Support`` / ``~/Library/Application Support`` / ``~/.local/share``) is computed by :func:`default\_os\_app\_home` but is ****NOT**** auto-selected here ? the launcher (P2) calls it and exports ``RALLY\_APP\_HOME`` explicitly. Keeping the resolver opt-in means importing this module has zero behavioural effect on the repo/dev/CI/test paths (no surprise writes into a real user profile during a test run).

stdlib-only (os, platform, pathlib) so it is safe to import at ``backend.main`` line 1 ? before the heavier pydantic-backed ``backend.config`` package ? keeping the early ``load\_dotenv`` path light.

```
"""
```

```
from __future__ import annotations
```

```
import os
```

```
import platform
```

```
from pathlib import Path
```

```
#: Environment variable the packaged launcher sets to pin the app-data home.
```

```
APP_HOME_ENV = "RALLY_APP_HOME"
```

```
#: OS-app-dir leaf used by :func:`default_os_app_home` (the product brand).
```

```
APP_DIR_NAME = "Khelsutra"
```

```
def app_home() -> Path:
```

```
    """The resolved application-data home.
```

```
    ``RALLY_APP_HOME`` (``~`` expanded) when set; otherwise ``Path(".")`` ? the CWD, byte-identical
```

```
    to the pre-P0a behaviour. Does NOT create the directory (callers create on write).
```

```
    """
```

```
    env = os.environ.get(APP_HOME_ENV)
```

```
    if env and env.strip():
```

```
        return Path(env.strip()).expanduser()
```

```
    return Path(".")
```

```
def app_path(*parts: str | os.PathLike[str]) -> Path:
```

```
    """A path under :func:`app_home` (e.g. ``app_path("output")``)."""
```

```
    return app_home().joinpath(*[str(p) for p in parts])
```

```
def default_config_path() -> Path:
```

```
    """``<app-home>/config.json`` ? the config file the loader reads and ``--setup`` writes.
```

```
    With ``RALLY_APP_HOME`` unset this is ``Path("config.json")`` (CWD-relative), unchanged.
```

```
    """
```

```
    return app_home() / "config.json"
```

```
def default_env_path() -> Path:
```

```
    """``<app-home>/.env`` ? the dotenv file the server prefers (e.g. ``GEMINI_API_KEY``)."""
```

```
    return app_home() / ".env"
```

```
def resolve_output_dir(output_dir: str) -> str:
```

```
    """Root a RELATIVE ``output_dir`` under :func:`app_home`; pass an ABSOLUTE one through.
```

```
    With ``RALLY_APP_HOME`` unset, ``resolve_output_dir("output") == "output"`` (CWD-relative), byte-identical to today; with it set, ``"output"`` becomes ``<app-home>/output`` so the
```


SQLite

DB and produced reels live alongside the rest of the app's state.

"""

```
if not output_dir:
    return str(app_home() / "output")
if os.path.isabs(output_dir):
    return output_dir
return str(app_home() / output_dir)
```

def default_os_app_home(app_name: str = APP_DIR_NAME) -> Path:

"""The OS-conventional per-user app-data dir (for the packaged LAUNCHER to export as
`RALLY\_APP\_HOME` ? NOT auto-selected by :func:`app\_home`).

* Windows ? ``%APPDATA%\<app\_name>`` (``~/AppData/Roaming/<app\_name>`` if unset)
* macOS ? ``~/Library/Application Support/<app\_name>``
* Linux/other ? ``\$XDG\_DATA\_HOME/<app\_name>`` (`` ~/.local/share/<app\_name>`` if unset)
"""

```
system = platform.system()
if system == "Windows":
    base = os.environ.get("APPDATA") or str(Path.home() / "AppData" / "Roaming")
elif system == "Darwin":
    base = str(Path.home() / "Library" / "Application Support")
else:
    base = os.environ.get("XDG_DATA_HOME") or str(Path.home() / ".local" / "share")
return Path(base) / app_name
```

76. user (2026-06-21T14:32:37.685Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0a\backend\utils\app_home.py (file state is current in your context ? no need to Read it back)

77. assistant (2026-06-21T14:32:40.917Z)

Now thread it through the loader, main, and doctor.

78. user (2026-06-21T14:32:45.375Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\config\loader.py has been updated successfully. (file state is current in your context ? no need to Read it back)

79. user (2026-06-21T14:32:52.360Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\config\loader.py has been updated successfully. (file state is current in your context ? no need to Read it back)

80. user (2026-06-21T14:32:56.446Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\config\loader.py has been updated successfully. (file state is current in your context ? no need to Read it back)

81. assistant (2026-06-21T14:32:59.614Z)

Now `main.py` ? the `.env` load, the static mount, and the output-dir resolution at server start + trim.

82. user (2026-06-21T14:33:02.701Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

83. user (2026-06-21T14:33:13.377Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

84. user (2026-06-21T14:33:28.069Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

85. user (2026-06-21T14:33:32.354Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

86. assistant (2026-06-21T14:33:35.065Z)

Now `doctor.py` ? unify its config write with the server's read.

87. user (2026-06-21T14:33:39.332Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\scripts\doctor.py has been updated successfully. (file state is current in your context ? no need to Read it back)

88. assistant (2026-06-21T14:33:51.721Z)

Now the tests.

89. assistant (2026-06-21T14:34:13.662Z)

```
"""Tests for the resolved application-data HOME (PACKAGING_PLAN.md P0a).
```

```
The load-bearing guarantee: with ``RALLY_APP_HOME`` UNSET, every resolver is byte-identical to the pre-P0a CWD-relative behaviour; with it SET, config/.env/output/DB all root under that one home.
"""
```

```
import os
from pathlib import Path
```

```
import pytest
```

```
from backend.utils import app_home as ah
from backend.config import loader
```

```
@pytest.fixture(autouse=True)
def _clear_app_home_env(monkeypatch):
    """Each test starts with RALLY_APP_HOME unset unless it sets it explicitly."""
    monkeypatch.delenv(ah.APP_HOME_ENV, raising=False)
```

---- default (env unset) is byte-identical to today ----

```
def test_app_home_defaults_to_cwd():
    assert ah.app_home() == Path(".")
```

```
def test_default_config_path_unset_is_cwd_config_json():
    # Pre-P0a this was the module constant Path("config.json").
    assert ah.default_config_path() == Path("config.json")
```

```
def test_default_env_path_unset_is_cwd_dotenv():
    assert ah.default_env_path() == Path(".env")
```

```
def test_resolve_output_dir_unset_is_identity_for_relative():
    assert ah.resolve_output_dir("output") == "output"
    assert ah.resolve_output_dir("custom") == "custom"
```

---- RALLY_APP_HOME set ? everything roots there ----

```
def test_app_home_honors_env(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    assert ah.app_home() == tmp_path
```

```
def test_app_home_expands_user(monkeypatch):
    monkeypatch.setenv(ah.APP_HOME_ENV, "~/khelsutra-data")
    assert ah.app_home() == Path.home() / "khelsutra-data"
```

```
def test_app_home_ignores_blank_env(monkeypatch):
    monkeypatch.setenv(ah.APP_HOME_ENV, " ")
    assert ah.app_home() == Path(".")
```

```
def test_config_and_env_paths_root_under_home(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    assert ah.default_config_path() == tmp_path / "config.json"
    assert ah.default_env_path() == tmp_path / ".env"
    assert ah.app_path("output", "x.db") == tmp_path / "output" / "x.db"
```

```
def test_resolve_output_dir_set(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    assert ah.resolve_output_dir("output") == str(tmp_path / "output")
```

```
def test_resolve_output_dir_absolute_passthrough(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    abs_dir = str(tmp_path / "elsewhere") if os.path.isabs(str(tmp_path)) else
os.path.abspath("elsewhere")
    assert ah.resolve_output_dir(abs_dir) == abs_dir
```

```
def test_resolve_output_dir_empty_falls_back_to_home_output(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    assert ah.resolve_output_dir("") == str(tmp_path / "output")
```

---- default_os_app_home: the OS-conventional dir the launcher exports (NOT auto-se

```
def test_default_os_app_home_windows(monkeypatch, tmp_path):
    monkeypatch.setattr(ah.platform, "system", lambda: "Windows")
    monkeypatch.setenv("APPDATA", str(tmp_path / "Roaming"))
    assert ah.default_os_app_home() == tmp_path / "Roaming" / "Khelsutra"
```

```
def test_default_os_app_home_macos(monkeypatch):
    monkeypatch.setattr(ah.platform, "system", lambda: "Darwin")
    assert ah.default_os_app_home() == Path.home() / "Library" / "Application Support" /
"Khelsutra"
```

```
def test_default_os_app_home_linux_xdg(monkeypatch, tmp_path):
    monkeypatch.setattr(ah.platform, "system", lambda: "Linux")
    monkeypatch.setenv("XDG_DATA_HOME", str(tmp_path / "share"))
```

```

assert ah.default_os_app_home() == tmp_path / "share" / "Khelsutra"

def test_default_os_app_home_linux_default(monkeypatch):
    monkeypatch.setattr(ah.platform, "system", lambda: "Linux")
    monkeypatch.delenv("XDG_DATA_HOME", raising=False)
    assert ah.default_os_app_home() == Path.home() / ".local" / "share" / "Khelsutra"

```

---- loader integration: load_config() roots on the app-home ----

```

def test_loader_reads_config_from_app_home(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    (tmp_path / "config.json").write_text('{"sport": "tennis"}', encoding="utf-8")
    cfg = loader.load_config() # no explicit path ? default_config_path() under the home
    assert cfg.sport == "tennis"

def test_loader_no_config_in_home_uses_defaults(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path)) # empty dir, no config.json
    cfg = loader.load_config()
    # Defaults load cleanly (no crash, no file) ? same contract as a missing CWD config.json.
    assert cfg is not None
    assert hasattr(cfg, "indexing")

def test_save_default_config_writes_under_app_home(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    written = loader.save_default_config() # no explicit path
    assert written == tmp_path / "config.json"
    assert written.exists()

def test_explicit_path_still_overrides_app_home(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path / "home"))
    explicit = tmp_path / "explicit.json"
    explicit.write_text('{"sport": "squash"}', encoding="utf-8")
    cfg = loader.load_config(str(explicit))
    assert cfg.sport == "squash"

```

90. user (2026-06-21T14:34:13.689Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0a\tests\test_app_home.py (file state is current in your context ? no need to Read it back)

91. assistant (2026-06-21T14:34:34.348Z)

Now run the targeted tests, but first verify the worktree's `backend` resolves to the worktree (not the editable install pointing at the main checkout).

92. user (2026-06-21T14:34:44.276Z)

```

=== which backend.main resolves ===
C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py
=== app_home importable ===
C:\Users\avidu\Projects\bhi-wt-p0a\backend\utils\app_home.py
home= .
cfg= config.json
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

93. assistant (2026-06-21T14:34:51.799Z)

Worktree code wins, and defaults are byte-identical (`app\_home()=.` , `config.json`). Now run the


```
..... [100%]
110 passed in 19.33s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

100. assistant (2026-06-21T14:35:41.182Z)

All green. Now the full suite + lint/type checks (the pre-LGTM gate).

101. user (2026-06-21T14:36:11.167Z)

```
=== RUFF (changed files) ===
All checks passed!
=== MYPY (changed modules) ===
Success: no issues found in 2 source files
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

102. assistant (2026-06-21T14:36:18.681Z)

Ruff + mypy clean. Now the full suite (the pre-LGTM gate ? whole suite, not just changed files).

103. user (2026-06-21T14:37:49.495Z)

```
..... [ 17%]
..... [ 23%]
..... [ 29%]
..... [ 35%]
..... [ 41%]
..... [ 46%]
..... [ 52%]
..... [ 58%]
..... [ 64%]
..... [ 70%]
.....SSSSS..... [ 76%]
..... [ 82%]
..... [ 87%]
..... [ 93%]
..... [ 99%]
.... [100%]
1221 passed, 7 skipped in 86.76s (0:01:26)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

104. assistant (2026-06-21T14:37:54.917Z)

Full suite green (1221 passed, 7 skipped ? the optional obsreport tests). Let me confirm the coverage floor and that my change holds it.

105. user (2026-06-21T14:37:56.712Z)

```
=== coverage floor in CI / pytest.ini ===
.github/workflows/ci.yml:97: run: python run_all_tests.py --cov=backend --cov=training
--cov-report=term --cov-fail-under=70
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

106. user (2026-06-21T14:39:50.167Z)

```
Full test suite: C:\Users\avidu\AppData\Local\Programs\Python\Python313\python.exe -m pytest
-q --cov=backend --cov=training --cov-report=term --cov-fail-under=70
tests/test_motion_segmenter.py::test_process_video_failure_on_unopenable_video
tests/test_motion_segmenter.py::test_process_video_failure_on_unopenable_video
tests/test_motion_segmenter.py::test_process_video_failure_on_unopenable_video
backend\config\loader.py 91 3 97%
```

```

backend\eval\gt_loader.py          55      0  100%
backend\utils\app_home.py         31      0  100%
TOTAL                             9623  1434   85%
Required test coverage of 70% reached. Total coverage: 85.10%
RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

107. assistant (2026-06-21T14:40:16.318Z)

Coverage ****85.10%**** (floor 70; no regression vs the ~85% baseline) ? `app\_home.py` 100%, `loader.py` 97%. Now the owner-preferred ****proof****: demonstrate the user's state (config + SQLite DB) actually lands under `RALLY\_APP\_HOME` when set, and in the CWD (unchanged) when not ? exercising the exact resolver + `Database` the server uses at `main.py:656-660`.

108. user (2026-06-21T14:40:26.565Z)

Exit code 1

=== CASE 1: RALLY_APP_HOME UNSET (today's behaviour ? CWD-relative) ===

```

app_home()      : .
config path     : config.json
output dir      : output (relative -> under CWD)
DB created under CWD?: True ->

```

C:\Users\avidu\AppData\Local\Temp\tmpmdt9oj5y\output\sports_indexer.db

Traceback (most recent call last):

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\shutil.py", line 631, in
_rmtree_unsafe
    os.rmdir(path)
    ~~~~~^~~~~~

```

PermissionError: [WinError 32] The process cannot access the file because it is being used by another process: 'C:\\Users\\avidu\\AppData\\Local\\Temp\\tmpmdt9oj5y'

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\tempfile.py", line 910, in
onexc
    _os.unlink(path)
    ~~~~~^~~~~~

```

PermissionError: [WinError 5] Access is denied:
'C:\\Users\\avidu\\AppData\\Local\\Temp\\tmpmdt9oj5y'

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

```

File "<stdin>", line 17, in <module>
File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\tempfile.py", line 951, in
__exit__
    self.cleanup()
    ~~~~~^~~~~~

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\tempfile.py", line 955, in cleanup

```

    self._rmtree(self.name, ignore_errors=self._ignore_cleanup_errors)
    ~~~~~^~~~~~

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\tempfile.py", line 935, in _rmtree

```

    _shutil.rmtree(name, onexc=onexc)
    ~~~~~^~~~~~

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\shutil.py", line 790, in rmtree

```

    return _rmtree_unsafe(path, onexc)

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\shutil.py", line 635, in _rmtree_unsafe

```

    onexc(os.rmdir, path, err)
    ~~~~~^~~~~~

```

File "C:\Users\avidu\AppData\Local\Programs\Python\Python313\Lib\tempfile.py", line 925, in


```
backend/main.py          | 29 ++++++-----
scripts/doctor.py        | 6 +++++-
3 files changed, 35 insertions(+), 11 deletions(-)
=== untracked new files ===
?? backend/utils/app_home.py
?? tests/test_app_home.py
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

113. assistant (2026-06-21T14:41:08.443Z)

Diff is tight (3 files +35/-11, 2 new files). Running a 4-lens adversarial review over the worktree before commit.

114. user (2026-06-21T14:41:32.898Z)

Workflow launched in background. Task ID: w00dr4axu
Summary: Adversarial review of the P0a app-data-home PR (behaviour-neutrality, import-safety, cross-platform paths, test gaps)
Transcript dir: C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\subagents\workflows\wf_502b2e43-89d
Script file: C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js
(Edit this file with Write/Edit and re-invoke Workflow with {scriptPath: "C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js"} to iterate without resending the script.)
Run ID: wf_502b2e43-89d
To resume after editing the script: Workflow({scriptPath: "C:\Users\avidu\claudu\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js", resumeFromRunId: "wf_502b2e43-89d"}) ? completed agents return cached results.

You will be notified when it completes. Use /workflows to watch live progress.

115. assistant (2026-06-21T14:41:48.719Z)

Script bug ? `await parallel(...).filter(...)` binds as `await (parallel(...).filter)`, calling `.filter` on a Promise. Need `(await parallel(...)).filter(...)`. Let me fix the script file and re-run.

116. user (2026-06-21T14:41:52.433Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

117. user (2026-06-21T14:41:58.160Z)

```
24      { key:'import-safety-ordering', prompt:`Read backend/main.py top +
backend/utils/app_home.py. ADVERSARIALLY check import-time safety: main.py:1-20 imports app_home
BEFORE the heavy backend.config ? does importing backend.utils.app_home pull anything heavy or
do filesystem I/O at import? Is load_dotenv ordering preserved (app-home .env if exists else
walk-up)? The static mount does os.makedirs on a PACKAGE-relative dir ? on an INSTALLED wheel
(site-packages, possibly read-only) does makedirs(exist_ok=True) raise? Is that a regression vs
the old CWD makedirs? Flag any import-time crash or I/O risk.` },
25      { key:'cross-platform-paths', prompt:`Read backend/utils/app_home.py. ADVERSARIALLY
hunt cross-platform/Windows path bugs: expanduser on RALLY_APP_HOME; os.path.isabs semantics for
a Windows drive-relative path like "C:foo" or a UNC path; str(Path(".")."output")=="output" on
Windows; default_os_app_home for Windows (%APPDATA%), macOS, Linux (XDG) ? are the leafs correct
and is APP_DIR_NAME 'Khelsutra' fine? Any case where resolve_output_dir produces a surprising
path? CRLF/encoding when save_default_config writes?` },
26      { key:'test-gaps-and-doctor', prompt:`Read tests/test_app_home.py + scripts/doctor.py
+ the diff. What behaviours are UNtested? (e.g. the main.py load_dotenv branch, the package-
relative static mount, resolve_output_dir at the server-start/trim sites, doctor writing to
```

```
default_config_path). Also: the doctor change moves the config write from REPO_ROOT to
default_config_path() (CWD when env unset) ? is the behaviour change for "doctor run from a repo
SUBDIR" acceptable/intended, and is it documented? Suggest the specific missing tests.` },
27     ]
28
29     const results = await parallel(lenses.map(l => () =>
30         agent(`${CTX}\n\n${FILES}\n\nLENS: ${l.key}\n\n${l.prompt}\n\nRead the actual files
under ${WT}. Return findings with severity blocker|major|minor|nit, each with file:line and a
concrete fix. Be harsh but precise; do not invent issues.` ,
31         { label:`review:${l.key}`, phase:'Review', schema: SCHEMA })
```

118. user (2026-06-21T14:42:02.566Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js has been updated successfully. (file state is current in your context ? no need to Read it back)

119. assistant (2026-06-21T14:42:08.866Z)

(Preview verification is N/A ? this is a CLI/server packaging change, not browser-observable.)
Let me fix the matching closing paren and re-run the review.

120. user (2026-06-21T14:42:11.057Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js has been updated successfully. (file state is current in your context ? no need to Read it back)

121. user (2026-06-21T14:42:18.567Z)

Workflow launched in background. Task ID: wrzaqbx01
Summary: Adversarial review of the P0a app-data-home PR (behaviour-neutrality, import-safety, cross-platform paths, test gaps)
Transcript dir: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\subagents\workflows\wf_0f586cf2-9eb
Script file: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js
(Edit this file with Write/Edit and re-invoke Workflow with {scriptPath: "C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js"} to iterate without resending the script.)
Run ID: wf_0f586cf2-9eb
To resume after editing the script: Workflow({scriptPath: "C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\workflows\scripts\p0a-app-home-review-wf_502b2e43-89d.js", resumeFromRunId: "wf_0f586cf2-9eb"}) ? completed agents return cached results.

You will be notified when it completes. Use /workflows to watch live progress.

122. assistant (2026-06-21T14:42:43.030Z)

P0a is implemented and verified green (full suite 1221 passed, ruff/mypy clean, coverage 85.10%, and an end-to-end proof that state follows `RALLY\_APP\_HOME` and is byte-identical to today when unset). The 4-lens adversarial review is running; I'll fold any findings, then commit and open the PR. Waiting on the review rather than polling.

While that runs, here's the **P0a change set** (in the isolated worktree, branched from the fresh `origin/master` `7835f34`):

| File | Change |

```
|---|---|
| `backend/utils/app_home.py` *(new)* | The resolver: `app_home()` = `$RALLY_APP_HOME` or
`Path(".")`; `default_config_path` / `default_env_path` / `resolve_output_dir`;
`default_os_app_home()` (per-OS dir, used later by the P2 launcher, not auto-selected) |
| `backend/config/loader.py` | `load_config`/`save_default_config` resolve via
`default_config_path`; `DEFAULT_CONFIG_PATH` kept as a back-compat alias |
| `backend/main.py` | `.env` prefers app-home then falls back to today's search; static mount is
package-relative (not CWD); `output_dir` (DB + reels) routed through `resolve_output_dir` at
server-start and trim |
| `scripts/doctor.py` | writes `config.json` to `default_config_path()` ? same place the server
reads (fixes the doctor?server disagreement footgun) |
| `tests/test_app_home.py` *(new)* | 22 tests: byte-identical-when-unset, env honoring, per-OS
dirs, loader integration |
```

The contract held: ****zero behaviour change when `RALLY\_APP\_HOME` is unset****, and with it set, config + DB + reels all root in the app home with nothing leaking to the launch directory.

123. user (2026-06-21T14:45:47.256Z)

```
<task-notification>
<task-id>wrzaqbx01</task-id>
<tool-use-id>toolu_01F3AJTs8eZoUfsezs6GbGDh</tool-use-id>
<output-file>C:\Users\avidu\AppData\Local\Temp\claudex-C--Users-avidu-Projects-badminton-
highlight-indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\tasks\wrzaqbx01.output</output-file>
<status>completed</status>
<summary>Dynamic workflow "Adversarial review of the P0a app-data-home PR (behaviour-neutrality,
import-safety, cross-platform paths, test gaps)" completed</summary>
<result>[{"findings":[{"severity":"major","title":"load_dotenv() now discovers a CWD .env that
pre-P0a ignored ? unset path is NOT byte-
identical","where":"backend/main.py:15-19","detail":"Pre-P0a was a bare load_dotenv(), which
calls python-dotenv's find_dotenv(usecwd=False). For a normal .py run that does NOT search the
CWD ? it starts the walk-up from the CALLER FILE's directory (backend/), so it finds a repo-root
.env but NOT a .env sitting in an unrelated CWD. P0a replaced this with: _env_path =
default_env_path() (== Path('.env') when RALLY_APP_HOME is unset); if _env_path.exists():
load_dotenv(_env_path). When CWD is not an ancestor of backend/ and a .env exists in that CWD,
the new code loads it while pre-P0a would not. I reproduced this directly: with a .env only in
the CWD (CWD outside the package tree), pre-P0a load_dotenv() left the var unset (None) while
the P0a branch loaded it ('cwd_value'). The reverse case (repo-root .env, no CWD .env) is
identical, and the canonical repo-root run (CWD==repo root, which IS an ancestor of backend/) is
identical ? but the contract says 'env unset => byte-identical', and for the CWD-.env case it
is measurably not. This is exactly the arbitrary-CWD scenario the PR targets (though in real
packaging the launcher sets RALLY_APP_HOME, so the unset branch wouldn't run). No test guards
this ? test_app_home.py has no load_dotenv coverage.", "suggested_fix":"Either (a) make the unset
path explicitly fall through to the historical search ? only prefer an app-home .env when
RALLY_APP_HOME is actually SET: `import os; _home=os.environ.get('RALLY_APP_HOME'); if _home and
(Path(_home).expanduser())/'.env').exists(): load_dotenv(...) else: load_dotenv()` ? so the unset
branch is the verbatim pre-P0a call; or (b) consciously accept the new 'CWD .env wins'
precedence and update the hard-contract wording + add a test documenting it. Given the stated
HARD CONTRACT, prefer (a)."}, {"severity":"minor","title":"resolve_output_dir normalizes relative
paths (./ , forward slashes, trailing slash) ? not verbatim like the old
args.output_dir","where":"backend/utils/app_home.py:77 (consumed at
backend/main.py:634,671)","detail":"Old code stored args.output_dir VERBATIM (out_dir =
args.output_dir; global_config.output.output_dir = args.output_dir). New code routes a relative
value through str(app_home() / output_dir) = str(Path('.') / v), which pathlib NORMALIZES.
Verified on Windows: './output'->'output', 'a/b'->'a\\b', 'output/'->'output' (the
default 'output' and 'out\\sub' are invariant). The on-disk effect is the same
(os.makedirs/os.path.join accept either separator), but the STORED string differs and that
string is observable: it is echoed by GET /api/config, embedded in the per-video report, and
used to build download paths. So for a non-default --output-dir like './myout' or a forward-
slash path on Windows the unset behaviour is not byte-identical. Common case (--output-dir
default 'output') is unaffected.", "suggested_fix":"If strict byte-identity for relative inputs
matters: when app_home()==Path('.') (unset), return output_dir unchanged ? e.g. `home =
app_home(); if str(home) == '.': return output_dir; return str(home / output_dir)` (absolute
branch unchanged). Otherwise, broaden test_resolve_output_dir_unset_is_identity_for_relative to
assert the normalized values for './output', 'a/b', 'output/' so the normalization is documented
```

```

intent rather than an accident."},{ "severity": "nit", "title": "Test 'identity for relative' only
covers normalization-invariant inputs, giving false confidence", "where": "tests/test_app_home.py:
37-39", "detail": "test_resolve_output_dir_unset_is_identity_for_relative asserts only
resolve_output_dir('output')== 'output' and ('custom')== 'custom' ? both single-segment names that
survive pathlib normalization unchanged. It therefore does NOT actually prove 'identity for
relative' inputs; './output', 'a/b' (Windows), and 'output/' all change (see the minor finding).
The test name overclaims.", "suggested_fix": "Either rename to
test_resolve_output_dir_unset_single_segment_identity, or add cases for './output', a multi-
segment path, and a trailing-slash path asserting the real (normalized) outputs so the contract
surface is honest."}, {"severity": "nit", "title": "DEFAULT_CONFIG_PATH alias is import-time frozen
and now mildly misleading; CODE_MAP.md doc is stale", "where": "backend/config/loader.py:35 and
docs/CODE_MAP.md:132", "detail": "DEFAULT_CONFIG_PATH = default_config_path() is evaluated ONCE at
import. If RALLY_APP_HOME is set after import, this constant diverges from the dynamic
default_config_path() used by load_config/save_default_config. Confirmed there are no runtime
consumers (grep shows only the def site + a doc line), so it is harmless TODAY, but it is a
latent footgun for a future reader who reaches for the 'constant'. Separately,
docs/CODE_MAP.md:132 still documents it as the literal `Path(\"config.json\")` (CWD-relative),
which is now stale ? it is the resolver's value.", "suggested_fix": "Drop DEFAULT_CONFIG_PATH
entirely (no external consumers per the comment), OR if kept for back-compat make it a lazy
accessor / clearly mark 'snapshot at import; use default_config_path()'. Update
docs/CODE_MAP.md:132 to reflect that the loader now resolves via
backend.utils.app_home.default_config_path() (RALLY_APP_HOME-aware; ==config.json when
unset)."}], "overall": "Verdict: the CORE resolver contract holds ? I empirically verified on BOTH
Windows and POSIX that with RALLY_APP_HOME unset, default_config_path()==Path('config.json'),
default_env_path()==Path('.env'), and resolve_output_dir('output')== 'output' (pathlib collapses
the Path('.') prefix, so the strings are byte-identical). The static-mount change to package-
relative _STATIC_DIR is sound and byte-identical when CWD==repo root, with no other readers of
the old 'backend/api/static' literal. However, the HARD CONTRACT ('env unset => byte-
identical to pre-P0a') is violated in two demonstrated ways: (1) MAJOR ? the new load_dotenv
branch loads a CWD .env that pre-P0a's find_dotenv(usecwd=False) frame-walk would have ignored
(reproduced: pre-P0a left the var None, P0a loaded it); the canonical repo-root run is
unaffected, but the contract is about the general unset case. (2) MINOR ? resolve_output_dir
normalizes relative paths (./x, forward slashes, trailing slash) whereas the old code stored
args.output_dir verbatim, so a non-default --output-dir yields a different (though functionally-
equivalent on disk) stored/echoed string. Plus two nits: the 'identity' test only exercises
normalization-invariant inputs (false confidence), and the import-frozen DEFAULT_CONFIG_PATH
alias + a stale CODE_MAP.md line. The full app_home test module passes (19/19). Recommend gating
the dotenv app-home preference behind an actually-SET RALLY_APP_HOME (verbatim fall-through when
unset) to restore strict byte-identity, and broadening the output-dir
test."}, {"findings": [{"severity": "minor", "title": "Import-time os.makedirs into the install tree
is unguarded (PermissionError not swallowed by
exist_ok=True)", "where": "backend/main.py:84", "detail": "_STATIC_DIR is now package-relative
(inside site-packages on a wheel) and os.makedirs(_STATIC_DIR, exist_ok=True) runs at IMPORT
time of backend.main. I built the wheel and confirmed the static assets DO ship
(backend/api/static/{index.html,app.js,style.css} are in the wheel), so on a normal install the
dir already exists and makedirs is a clean no-op early-return -- NOT a regression for the
product as shipped today. The risk is forward-looking: exist_ok=True suppresses ONLY
FileExistsError, never PermissionError (verified: a missing leaf under a read-only parent still
raises). The old code created ./backend/api/static under the CWD, which is always user-writable.
If a future packaging step (P1a explicitly reworks static bundling) ever ships the package
WITHOUT that dir, and site-packages is read-only (system/root-installed Python on
... (truncated 23474 chars, full result in C:\Users\avidu\AppData\Local\Temp\claude\C--Users-
avidu-Projects-badminton-highlight-
indexer\c7679509-ee2b-4443-b240-0b3b1b1a7291\tasks\wrzaqbx01.output)</result>
<usage><agent_count>4</agent_count><subagent_tokens>314180</subagent_tokens><tool_uses>97</tool_
uses><duration_ms>208645</duration_ms></usage>
</task-notification>

```

124. assistant (2026-06-21T14:46:34.111Z)

Strong review ? one MAJOR byte-identity violation and several worth fixing. Let me read the full output for the remaining lenses.

125. user (2026-06-21T14:46:38.094Z)

[MAJOR] load_dotenv() now discovers a CWD .env that pre-P0a ignored ? unset path is NOT byte-identical

where: backend/main.py:15-19

fix: Either (a) make the unset path explicitly fall through to the historical search ? only prefer an app-home .env when RALLY_APP_HOME is actually SET: ``import os; _home=os.environ.get('RALLY_APP_HOME'); if _home and (Path(_home).expanduser()/.env').exists(): load_dotenv(...) else: load_dotenv()`` ? so the unset branch is the verbatim pre-P0a call; or (b) consciously accept the new 'CWD .env wins' precedence and update the hard-contract wording + add a test documenting it. Given the stated HARD CONTRACT, prefer (a).

[MAJOR] resolve_output_dir() violates the "byte-identical when unset" HARD CONTRACT for relative paths containing separators / dot / trailing slash

where: backend/utils/app_home.py:77 (resolve_output_dir relative branch); contract asserted at backend/utils/app_home.py:9-11 and 66-72

fix: Short-circuit the unset case so it is a true identity. e.g. at the top of resolve_output_dir: ``if not os.environ.get(APP_HOME_ENV, '').strip(): return output_dir or 'output'`` (return the raw string untouched when no app-home is pinned). Or more generally, when `app_home()==Path('.')`, return output_dir verbatim. Then add unset-mode tests covering ``out/clips`, `./output`, `output/`` asserting identity to the input.

[MAJOR] load_dotenv app-home branch (main.py:15-19) is entirely untested ? the core P0a behaviour change is unverified

where: backend/main.py:15-19; tests/test_app_home.py (no coverage)

fix: Extract the env-resolution into a pure helper, e.g. in app_home.py: ``def env_file_to_load() -> Path | None:`` returning default_env_path() if it exists else None (None => caller uses bare load_dotenv() walk-up). Then test all four cases with monkeypatch+tmp_path: (1) env unset, no CWD .env => None (walk-up preserved); (2) env unset, CWD .env present => Path('.env'); (3) RALLY_APP_HOME set with .env in home => that path; (4) RALLY_APP_HOME set, home has no .env, CWD has one => home path still chosen / None-then-fallback per intended precedence. Assert the home .env wins over a CWD .env when the env is set.

[MAJOR] Package-relative static mount + serve_index (main.py:83-93) untested ? was a CWD-relative path, now package-relative

where: backend/main.py:83-93; tests/test_app_home.py (no coverage)

fix: Add a unit test asserting ``backend.main._STATIC_DIR == os.path.join(os.path.dirname(backend.main.__file__), 'api', 'static')`` (i.e. package-relative, independent of `os.getcwd()`). Add a TestClient test for GET ``/'``: monkeypatch ``_STATIC_DIR`` to a tmp_path, write an index.html, assert it is served; then assert the ``<h3>...Running...`` fallback is returned when index.html is absent. Optionally assert `os.makedirs` targeted ``_STATIC_DIR`` (not a CWD-relative path) by running from a `chdir'd` tmp dir and confirming no `./backend/api/static` is created.

[MAJOR] resolve_output_dir at the two main() wiring sites (server-start db path + --trim) is untested end-to-end

where: backend/main.py:634 (trim) and main.py:671-674 (server-start output_dir + db_path); tests/test_app_home.py only unit-tests resolve_output_dir() in isolation

fix: Add a subprocess test (mirroring test_server_entrypoint.py) that sets env `RALLY_APP_HOME=tmp_path`, runs ``python -m backend.main --server --output-dir output`` from a DIFFERENT cwd, and asserts the DB file appears at `tmp_path/output/<db_name>` and NOT at `cwd/output`. Add the symmetric env-unset case asserting CWD/output is used. For --trim, a faster unit test: monkeypatch VideoCompiler + compute and assert resolve_output_dir was applied to the computed out_dir under a set RALLY_APP_HOME. At minimum, set RALLY_APP_HOME in the existing test_server_entrypoint subprocess and assert the out_dir lands under the home.

[MAJOR] doctor.init_default_config writing to default_config_path() is untested, and the run-from-subdir behaviour change is undocumented for users

where: scripts/doctor.py:61; docs (no SETUP_NEW_MACHINE/PACKAGING note); tests (none for doctor)

fix: (a) Add tests/test_doctor.py: with `RALLY_APP_HOME=tmp_path`, call `init_default_config()` and assert `tmp_path/config.json` is created with valid AppConfig JSON; assert the skip-path prints/returns when the file already exists; assert the import-failure branch (monkeypatch the import to raise) is a no-op. (b) Fix the now-false docstring at `scripts/doctor.py:6-9` ? it claims the script 'operates on the repo root (config.json)'; `config.json` now follows `default_config_path()` (`RALLY_APP_HOME` else CWD), only requirements.txt is still `REPO_ROOT`-relative. (c) Note the run-from-subdir CWD behaviour in `PACKAGING_PLAN` P0a (and/or

SETUP_NEW_MACHINE) so the change is recorded as intended.

[MINOR] resolve_output_dir normalizes relative paths (./ , forward slashes, trailing slash) ? not verbatim like the old args.output_dir

where: backend/utils/app_home.py:77 (consumed at backend/main.py:634,671)

fix: If strict byte-identity for relative inputs matters: when app_home()==Path('.') (unset), return output_dir unchanged ? e.g. `home = app\_home(); if str(home) == '.': return output\_dir; return str(home / output\_dir)` (absolute branch unchanged). Otherwise, broaden test_resolve_output_dir_unset_is_identity_for_relative to assert the normalized values for './output', 'a/b', 'output/' so the normalization is documented intent rather than an accident.

[MINOR] Import-time os.makedirs into the install tree is unguarded (PermissionError not swallowed by exist_ok=True)

where: backend/main.py:84

fix: Drop the import-time makedirs and just mount the shipped dir, or guard it: `try:\n os.makedirs(\_STATIC\_DIR, exist\_ok=True)\nexcept OSError:\n pass # shipped read-only data dir; never created at runtime`. StaticFiles(directory=...) tolerates a present dir without needing creation.

[MINOR] load_dotenv ordering is not strictly byte-identical: CWD .env now wins over the historical frame-relative walk-up

where: backend/main.py:15-19

fix: Either accept and note it in the PR (the dev/CI path is unaffected), or, to be strictly byte-identical when env is unset, only prefer the app-home .env when RALLY_APP_HOME is actually set: `from backend.utils.app\_home import app\_home, APP\_HOME\_ENV; if os.environ.get(APP\_HOME\_ENV) and default\_env\_path().exists(): load\_dotenv(default\_env\_path()) else: load\_dotenv()` so the unset path keeps the exact historical find_dotenv() walk-up.

[MINOR] Windows drive-relative / driveless-rooted output_dir is silently mis-rerooted when RALLY_APP_HOME is set

where: backend/utils/app_home.py:75-77

fix: Use a Windows-aware absoluteness/anchor check rather than bare os.path.isabs. e.g. treat `Path(output\_dir)` as pass-through when it `.is\_absolute()` OR has a drive (`PureWindowsPath(output\_dir).drive`) OR is anchored; alternatively normalize via `(app\_home() / output\_dir).resolve()`-style anchoring and document that drive-relative inputs are rejected. At minimum add a comment + test pinning the chosen behaviour for `C:foo` and `\\foo`.

[MINOR] No test asserts the env-unset 'byte-identical' contract for the COMPOSED main() flow ? only the resolver primitives

where: tests/test_app_home.py:24-39

fix: Once the per-site tests above exist, add an explicit 'env unset == pre-P0a' assertion at each: load chooses bare load_dotenv() walk-up; resolve_output_dir('output')== 'output' at the server-start site so db_path is os.path.join('output', db_name); doctor writes ./config.json. Group them so the contract is visibly enforced at the call sites, not just the helpers.

[NIT] Test 'identity for relative' only covers normalization-invariant inputs, giving false confidence

where: tests/test_app_home.py:37-39

fix: Either rename to test_resolve_output_dir_unset_single_segment_identity, or add cases for './output', a multi-segment path, and a trailing-slash path asserting the real (normalized) outputs so the contract surface is honest.

[NIT] DEFAULT_CONFIG_PATH alias is import-time frozen and now mildly misleading; CODE_MAP.md doc is stale

where: backend/config/loader.py:35 and docs/CODE_MAP.md:132

fix: Drop DEFAULT_CONFIG_PATH entirely (no external consumers per the comment), OR if kept for back-compat make it a lazy accessor / clearly mark 'snapshot at import; use default_config_path()'. Update docs/CODE_MAP.md:132 to reflect that the loader now resolves via backend.utils.app_home.default_config_path() (RALLY_APP_HOME-aware; ==config.json when unset).

[NIT] DEFAULT_CONFIG_PATH is import-time-frozen and has no code consumers -- a latent footgun

where: backend/config/loader.py:35

fix: Either delete the constant (no consumers) and update the docs/CODE_MAP.md line, or make it lazy/explicitly dynamic, e.g. drop the module-level binding and expose a function, so it cannot capture a stale RALLY_APP_HOME. If kept for back-compat, add a comment that it is frozen-

at-import and callers MUST use default_config_path() for the live value.

[NIT] save_default_config writes CRLF-terminated JSON on Windows (platform-divergent output bytes)

where: backend/config/loader.py:210 (write_text in text mode)

fix: If byte-stable output across platforms is wanted, write with an explicit newline: ``cfg_path.open("w", encoding="utf-8", newline="\n").write(...)``. Otherwise leave as-is ? low priority.

[NIT] DEFAULT_CONFIG_PATH back-compat alias is frozen at import time and never re-read ? keep it, but a test would prevent future misuse

where: backend/config/loader.py:35

fix: Either drop the alias entirely (the diff comment says there are no external consumers), or convert it to a property-like accessor / module `__getattr__` that calls `default_config_path()` on access. If kept, add a one-line test/comment that `DEFAULT_CONFIG_PATH` is import-time-frozen and the live helper must be used, so no one wires dynamic behaviour onto it.

126. assistant (2026-06-21T14:47:45.690Z)

The review confirmed the core resolver but caught two real byte-identity gaps (``load_dotenv`` CWD-``.env`` precedence; ``resolve_output_dir`` normalizing relative paths) plus untested wiring. All are legitimate ? I'll fix the two contract violations, make ``.env`` resolution a testable helper, add the missing tests (including codifying my end-to-end proof as a subprocess test), guard the import-time ``mkdirs``, drop the frozen ``DEFAULT_CONFIG_PATH``, and fix the stale doctor docstring + `CODE_MAP` line. Let me read the entrypoint test pattern and the `CODE_MAP` line first.

127. user (2026-06-21T14:47:46.657Z)

```
1      """Regression guard for the `python -m backend.main` double-module job-hang.
2
3      `python -m backend.main` runs main.py as `__main__`, so `main()` configures `__main__`'s
globals
4      (`db_instance`/`global_config`). But `job_worker`/`uploads` resolve shared state via
5      `import backend.main` ? a SEPARATE module object whose `db_instance` stays `None`. The
async-job
6      drain then crashes silently on `None.claim_due_job()` and every submitted job hangs in
`queued`
7      forever. The in-process tests (TestClient + monkeypatched `db_instance` + an inline
executor)
8      cannot see this ? only the real subprocess does. This spawns the DOCUMENTED invocation
and asserts
9      a submitted job actually LEAVES `queued`.
10     """
11     import json
12     import os
13     import socket
14     import subprocess
15     import sys
16     import time
17     import urllib.error
18     import urllib.request
19
20     import pytest
21
22     REPO_ROOT = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
23
24
25     def _free_port() -> int:
26         s = socket.socket()
27         s.bind(("127.0.0.1", 0))
28         port = int(s.getsockname()[1])
29         s.close()
30         return port
31
```

```

32
33     def _get(url: str, timeout: float = 5):
34         with urllib.request.urlopen(url, timeout=timeout) as r:
35             return json.loads(r.read().decode())
36
37
38     def _post(url: str, payload: dict, timeout: float = 15):
39         req = urllib.request.Request(
40             url, data=json.dumps(payload).encode(), headers={"content-type":
"application/json"})
41         with urllib.request.urlopen(req, timeout=timeout) as r:
42             return json.loads(r.read().decode())
43
44
45     def test_python_m_backend_main_drains_submitted_jobs(tmp_path):
46         """Spawning the engine the documented way (`python -m backend.main --server`) must
actually
47         DRAIN a submitted job ? it must leave `queued`, never hang. (A dummy file fails
fast; the
48         verdict is irrelevant ? what matters is that the worker claimed and ran it.)"""
49         port = _free_port()
50         out_dir = tmp_path / "out"
51         video = tmp_path / "dummy.mp4"
52         video.write_bytes(b"not-a-real-video") # exists -> the drain claims it; processing
fails fast
53         log = open(tmp_path / "server.log", "wb")
54
55         proc = subprocess.Popen(
56             [sys.executable, "-m", "backend.main", "--server", "--port", str(port),
57              "--output-dir", str(out_dir)],
58             cwd=REPO_ROOT, stdout=log, stderr=subprocess.STDOUT,
59         )
60         base = f"http://127.0.0.1:{port}"
61         try:
62             deadline = time.time() + 90 # cold import of torch/cv2 can be slow on CI
63             while time.time() < deadline:
64                 if proc.poll() is not None:
65                     pytest.fail(f"server exited early (code {proc.returncode}); see
{tmp_path/'server.log'}")
66                 try:
67                     if _get(f"{base}/api/health").get("status") == "ok":
68                         break
69                     except (urllib.error.URLError, ConnectionError, OSError):
70                         time.sleep(0.5)
71             else:
72                 pytest.fail("server never became healthy")
73
74             job_id = _post(f"{base}/api/process",
75                           {"video_path": str(video), "segmenter": "motion"})["job_id"]
76
77             status = "queued"
78             deadline = time.time() + 45
79             while time.time() < deadline:
80                 status = _get(f"{base}/api/jobs/{job_id}")["status"]
81                 if status != "queued":
82                     break
83                 time.sleep(1)
84
85             assert status != "queued", (
86                 "job stayed 'queued' ? the async drain never ran under `python -m
backend.main` "
87                 "(the __main__/backend.main double-module entrypoint bug regressed).")
88         )
89         finally:
90             proc.terminate()

```



```
91         try:
92             proc.wait(timeout=10)
93         except subprocess.TimeoutExpired:
94             proc.kill()
95         log.close()
96
```

128. assistant (2026-06-21T14:49:09.225Z)

The entrypoint test gives me a clean subprocess pattern (and I'll make my version `PYTHONPATH`-pin to this tree so it tests the worktree code, self-validating via DB location). Implementing all fixes now. First, `app\_home.py` ? verbatim-when-unset for `resolve\_output\_dir`, plus a testable `env\_file\_to\_load()` helper.

129. user (2026-06-21T14:49:19.325Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\utils\app_home.py has been updated successfully. (file state is current in your context ? no need to Read it back)

130. assistant (2026-06-21T14:49:23.124Z)

Now `main.py` ? use the `env\_file\_to\_load()` helper and guard the import-time `makedirs`.

131. user (2026-06-21T14:49:27.636Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

132. user (2026-06-21T14:49:33.729Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

133. assistant (2026-06-21T14:49:42.544Z)

Now drop the frozen `DEFAULT\_CONFIG\_PATH` (zero consumers ? the use-sites already call `default\_config\_path()`).

134. user (2026-06-21T14:49:46.412Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\backend\config\loader.py has been updated successfully. (file state is current in your context ? no need to Read it back)

135. assistant (2026-06-21T14:49:53.470Z)

Fix the now-stale doctor docstring.

136. user (2026-06-21T14:49:56.917Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\scripts\doctor.py has been updated successfully. (file state is current in your context ? no need to Read it back)

137. assistant (2026-06-21T14:50:01.221Z)

Now check CODE_MAP.md for the stale `DEFAULT\_CONFIG\_PATH` reference.

138. user (2026-06-21T14:50:02.462Z)

53:typed config load (built-in defaults -> config.json overrides; absent/parse-fail -> all-defaults, never fatal)

```
127: - `::AppConfig` ? root. Scalars `sport='badminton'`, `ai_provider='gemini'`,
    `ai_model='gemini-2.5-flash'`; sections `validation`, `indexing`, `output`, `stitching`
    (`default_factory`). `model_config = {extra:'allow', validate_assignment:True}`. ?
```

```

`extra='allow'` -> unknown/typo'd config.json keys silently kept (legacy tolerance).
129: - `::IndexingConfig` ? segmenter/detector knobs (`motion_threshold=0.08`,
`min_segment_duration=3.0`, `dead_time_merge_gap=2.0`, `chunk_duration=90.0`,
`chunk_overlap=10.0`, `yolo_velocity_threshold=0.15`, `yolo_frame_skip=3`,
`ai_handoff_padding=2.0`, `ai_max_workers=5`, `log/store_yolo_candidates=True`, nested
`tracknet`). ? `default_segmenter='yolo_hybrid'` HERE but config.json overrides to `gemini`. ?
`::IndexingConfig.segmenter_must_be_registered` `@field_validator` is a **no-op pass-through** ?
a typo'd segmenter validates fine, only fails later at `segmenter_registry.get(...)` .
132: - `::DEFAULT_CONFIG_PATH = Path("config.json")` (? CWD-relative, NOT repo-root).
`::get_built_in_defaults` = `AppConfig().model_dump(...)` .
133: - `::load_config` ? precedence built-in-defaults -> file overrides. ? **shallow top-level
merge** (`{**defaults, **file_data}`): a section in config.json REPLACES the whole defaults dict
for that section (Pydantic then re-applies field defaults for missing leaves).
Missing/unreadable/parse-error file -> fully-defaulted AppConfig + `[CONFIG WARNING]` (non-
fatal).
134: - `::save_default_config` ? writes fresh `config.json`, **overwrites** if present (used by
scripts/doctor.py/--setup). `::get_default_config` ? ? transitional plain-dict alias.
135:- `@config.json` ? on-disk config; mirrors `AppConfig`. ?
**`indexing.default_segmenter='gemini'`** diverges from model default `yolo_hybrid` ? file
wins at runtime.
136:- `@scripts/doctor.py` (formerly root `setup.py`) ? `::init_default_config` ? delegates to
`backend.config.save_default_config` (single-source defaults; skips if config.json exists).
`::run_pip_install` GPU-aware (`nvidia-smi` -> cu124 else cpu). `::run_diagnostics` (Python?3.8,
ffmpeg/ffprobe on PATH, optional torch+CUDA). Run via `python scripts/doctor.py` or `indexer
--setup`.
190:- `@backend/pipeline/stitcher/compiler.py::VideoCompiler` ? `compile_highlights(raw,
segments:[(start,end)], out_name)`. Per-clip re-encode -> `-f concat -c copy`. `::parse_time`
(mirrors gemini.parse_time; ? unparseable -> 0.0), `::get_video_duration`. ? if duration probe
fails (0.0) NO boundary validation; `begin_padding`/`end_padding` ctor defaults 0.5/1.0 but
every live call site (`/api/compile`, CLI trim, `scripts/run_process_and_stitch.py`) threads
`config.json` stitching.{begin_padding,end_padding} (typed default 0.5/0.5); temp clips in a
TemporaryDirectory under output_dir; raises ValueError/FileNotFoundError/RuntimeError; uses
`print()`. ? two-stage codec contract: clips re-encoded with IDENTICAL settings precisely so
final concat can stream-copy ? change a preset and concat may break. ? optional **branding
watermark** (`stitching.watermark`, default-on): `::watermark_filter` builds a `-vf` (drawtext
via `textfile=` + optional logo `movie/overlay`) applied to the per-clip re-encode only (concat
stays `-c copy`); a failed watermark encode retries the clip WITHOUT it (never nukes the reel,
and logs the ffmpeg stderr reason). Pass `watermark=` to `VideoCompiler(...)` (a
`WatermarkConfig` model or dict; `None` = off). Default font = bundled Apache-2.0
`::BUNDLED_FONT` (`backend/assets/fonts/RobotoSlab-Regular.ttf`) so it renders without
fontconfig; `font_path` overrides.
284:- **Segmenter default:** single source of truth = model
`IndexingConfig.default_segmenter='yolo_hybrid'`; `config.json` may override at runtime (ships
`gemini`). run_*.py read the model value (no literals); `main.py` keeps the defensive `motion`
getattr fallback noted above.
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

139. user (2026-06-21T14:50:08.137Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

140. user (2026-06-21T14:50:12.329Z)

```

131 - `@backend/config/loader.py` ? load/save with defaults-merge.
132 - `::DEFAULT_CONFIG_PATH = Path("config.json")` (? CWD-relative, NOT repo-root).
`::get_built_in_defaults` = `AppConfig().model_dump(...)` .

```

141. user (2026-06-21T14:50:18.235Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\docs\CODE_MAP.md has been updated successfully.
(file state is current in your context ? no need to Read it back)

142. assistant (2026-06-21T14:50:36.065Z)

Now update the tests: make the identity test prove true verbatim behavior, and add coverage for `env\_file\_to\_load`, the package-relative static dir, and drive passthrough.

143. user (2026-06-21T14:50:40.876Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\tests\test_app_home.py has been updated successfully. (file state is current in your context ? no need to Read it back)

144. user (2026-06-21T14:50:53.715Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\tests\test_app_home.py has been updated successfully. (file state is current in your context ? no need to Read it back)

145. assistant (2026-06-21T14:51:01.400Z)

Now a doctor test and the end-to-end server-wiring subprocess test. Let me check how `scripts/` modules are imported in existing tests.

146. user (2026-06-21T14:51:04.373Z)

```
=== scripts/__init__.py? ===
ls: cannot access 'scripts/__init__.py': No such file or directory
=== existing tests importing scripts.* or doctor ===
tests/test_nightly_regression.py
tests/__pycache__/test_nightly_regression.cpython-313-pytest-9.0.3.pyc
=== conf test sys.path tweaks ===
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

147. user (2026-06-21T14:51:08.190Z)

```
1:"""Tests for scripts/nightly_regression.py ? the pure summary/comparison/report core.
9:import importlib.util
10:import os
11:import sys
13:import pytest
15:# scripts/ isn't a package ? load the module by path. Register in sys.modules BEFORE exec so
16:# dataclasses can resolve the `from __future__ import annotations` string annotations.
18:    "scripts", "nightly_regression.py")
19:_spec = importlib.util.spec_from_file_location("nightly_regression", _SCRIPT)
20:nr = importlib.util.module_from_spec(_spec)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

148. user (2026-06-21T14:51:12.441Z)

```
14
15    # scripts/ isn't a package ? load the module by path. Register in sys.modules BEFORE
exec so
16    # dataclasses can resolve the `from __future__ import annotations` string annotations.
17    _SCRIPT = os.path.join(os.path.dirname(os.path.abspath(__file__)),
18        "scripts", "nightly_regression.py")
19    _spec = importlib.util.spec_from_file_location("nightly_regression", _SCRIPT)
20    nr = importlib.util.module_from_spec(_spec)
21    sys.modules["nightly_regression"] = nr
22    _spec.loader.exec_module(nr)
23
24
25    # ----- #
```

149. assistant (2026-06-21T14:51:27.071Z)

```
"""Tests for scripts/doctor.py::init_default_config (PACKAGING_PLAN.md P0a).
```

doctor's --setup config write now follows ``backend.utils.app\_home.default\_config\_path()`` ? the SAME resolver the server READS ? instead of the repo root, so the two can't disagree once the

```
app
is launched from an arbitrary CWD.
"""
```

```
import importlib.util
import os
import sys
```

```
import pytest
```

```
from backend.utils import app_home as ah
```

scripts/ isn't a package ? load the module by path (mirrors tests/test_nightly_regression.py)

```
_SCRIPT = os.path.join(os.path.dirname(os.path.dirname(os.path.abspath(__file__))),
                        "scripts", "doctor.py")
_spec = importlib.util.spec_from_file_location("doctor", _SCRIPT)
doctor = importlib.util.module_from_spec(_spec)
sys.modules["doctor"] = doctor
_spec.loader.exec_module(doctor)
```

```
@pytest.fixture(autouse=True)
```

```
def _clear_app_home_env(monkeypatch):
    monkeypatch.delenv(ah.APP_HOME_ENV, raising=False)
```

```
def test_init_writes_config_under_app_home(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    doctor.init_default_config()
    written = tmp_path / "config.json"
    assert written.exists(), "config.json must land under RALLY_APP_HOME (the server's read
path)"
    # And it is the SAME path the loader resolves ? the footgun this PR closes.
    assert written == ah.default_config_path()
```

```
def test_init_skips_when_config_already_exists(monkeypatch, tmp_path):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))
    existing = tmp_path / "config.json"
    existing.write_text('{"sport": "tennis"}', encoding="utf-8")
    doctor.init_default_config() # must NOT overwrite
    assert existing.read_text(encoding="utf-8") == '{"sport": "tennis"}'
```

```
def test_init_unset_writes_to_cwd(monkeypatch, tmp_path):
    # env unset ? CWD/config.json (the server, run from the same CWD, reads exactly this).
    monkeypatch.chdir(tmp_path)
    doctor.init_default_config()
    assert (tmp_path / "config.json").exists()
```

```
def test_init_is_noop_when_backend_import_fails(monkeypatch, tmp_path, capsys):
    monkeypatch.setenv(ah.APP_HOME_ENV, str(tmp_path))

    real_import = __builtins__["__import__"] if isinstance(__builtins__, dict) else
__builtins__.__import__
```

```
def _boom(name, *args, **kwargs):
    if name == "backend.config" or name.startswith("backend.config"):
        raise ImportError("simulated missing backend.config")
    return real_import(name, *args, **kwargs)
```

```
monkeypatch.setattr("builtins.__import__", _boom)
doctor.init_default_config() # must degrade to a no-op, never raise
```

```
assert not (tmp_path / "config.json").exists()
assert "Falling back to no-op" in capsys.readouterr().out
```

150. user (2026-06-21T14:51:27.099Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0a\tests\test_doctor.py (file state is current in your context ? no need to Read it back)

151. assistant (2026-06-21T14:52:00.509Z)

Now the end-to-end server-wiring test ? codifying my proof as a subprocess test that's `PYTHONPATH`-pinned to this tree (self-validating: the DB only lands under the home if the worktree's `resolve\_output\_dir` wiring ran).

152. user (2026-06-21T14:52:09.710Z)

The file C:\Users\avidu\Projects\bhi-wt-p0a\tests\test_server_entrypoint.py has been updated successfully. (file state is current in your context ? no need to Read it back)

153. assistant (2026-06-21T14:52:12.889Z)

Now run the new/changed tests to confirm the review-driven fixes hold.

154. user (2026-06-21T14:52:35.278Z)

```
..... [100%]
31 passed in 19.20s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

155. assistant (2026-06-21T14:52:39.999Z)

All 31 green (incl. the subprocess end-to-end and doctor tests). Now lint/type the full change set, then the whole suite with coverage.

156. user (2026-06-21T14:52:44.690Z)

```
=== RUFF ===
All checks passed!
=== MYPY ===
Success: no issues found in 2 source files
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

157. user (2026-06-21T14:54:44.964Z)

```
backend\config\loader.py          90      3    97%
backend\eval\gt_loader.py         55      0   100%
backend\utils\app_home.py         40      0   100%
TOTAL                             9631   1435    85%
Required test coverage of 70% reached. Total coverage: 85.10%
1232 passed, 7 skipped, 7 warnings in 112.06s (0:01:52)
  RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

158. assistant (2026-06-21T14:55:09.480Z)

Full suite **1232 passed / 7 skipped**, coverage **85.10%**, ruff+mypy clean, all review findings folded. Now commit (explicit paths only, per the shared-checkout protocol) and open the PR. First, re-verify the branch state.

159. user (2026-06-21T14:55:13.732Z)

```
=== branch + HEAD ===
feat/packaging-p0a-app-home
```

```
7835f34 feat(gcp): always-install Ops Agent on GPU VMs + a cloud-GPU dry-run guide (#291)
=== origin/master now ===
2afbadf fix(dryrun): install the watermark font in bootstrap + track the UI-trigger asks (#292)
=== my worktree changes (must be ONLY P0a files) ===
M backend/config/loader.py
M backend/main.py
M docs/CODE_MAP.md
M scripts/doctor.py
M tests/test_server_entrypoint.py
?? backend/utils/app_home.py
?? tests/test_app_home.py
?? tests/test_doctor.py
=== origin/master..HEAD (should be EMPTY pre-commit) ===
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

160. assistant (2026-06-21T14:55:24.179Z)

Clean ? exactly the 8 P0a files, no strays. (`origin/master` advanced to `2afbadf` #292; my base `7835f34` is its ancestor, so the PR diff is just mine.) Staging explicit paths and committing.

161. user (2026-06-21T14:55:36.339Z)

```
warning: in the working copy of 'backend/utils/app_home.py', LF will be replaced by CRLF the
next time Git touches it
warning: in the working copy of 'tests/test_app_home.py', LF will be replaced by CRLF the next
time Git touches it
warning: in the working copy of 'tests/test_doctor.py', LF will be replaced by CRLF the next
time Git touches it
=== staged ===
M backend/config/loader.py
M backend/main.py
A backend/utils/app_home.py
M docs/CODE_MAP.md
M scripts/doctor.py
A tests/test_app_home.py
A tests/test_doctor.py
M tests/test_server_entrypoint.py
=== committed ===
7c81c4d feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for config/.env/DB/output
(default-OFF)
=== origin/master..HEAD (should be ONLY my commit) ===
7c81c4d feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for config/.env/DB/output
(default-OFF)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

162. assistant (2026-06-21T14:55:41.858Z)

Clean commit (only my commit in `origin/master..HEAD`). Pushing and opening the PR.

163. user (2026-06-21T14:55:44.868Z)

```
remote:
remote: Create a pull request for 'feat/packaging-p0a-app-home' on GitHub by visiting:
remote:      https://github.com/Khelsutra/badminton-highlight-
indexer/pull/new/feat/packaging-p0a-app-home
remote:
branch 'feat/packaging-p0a-app-home' set up to track 'origin/feat/packaging-p0a-app-home'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch]      feat/packaging-p0a-app-home -> feat/packaging-p0a-app-home
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

164. user (2026-06-21T14:56:05.428Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/293>

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

165. assistant (2026-06-21T14:56:32.808Z)

**PR #293 opened: [feat(pkg): P0a ? resolved app-data
HOME](https://github.com/Khelsutra/badminton-highlight-indexer/pull/293).** Let me check CI
status.

166. user (2026-06-21T14:56:55.094Z)

Mypy (lenient baseline) pass 32s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326718
Ruff lint pass 8s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326682
Syntax sweep + import smoke (no GPU stack) pass 21s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326815
Full test suite (offline, mocked) pending 0 https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326676
Golden fixtures cross-OS (ubuntu-latest) pass 23s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326771
Golden fixtures cross-OS (windows-latest) pending 0
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326779
Golden fixtures cross-OS (macos-latest) pass 16s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326747
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

167. assistant (2026-06-21T14:57:04.312Z)

Most checks green; the full suite + Windows golden fixtures are still running. Let me wait for them to finish.

168. user (2026-06-21T14:58:29.232Z)

Golden fixtures cross-OS (ubuntu-latest) pass 23s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326771
Golden fixtures cross-OS (windows-latest) pass 1m5s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326779
Mypy (lenient baseline) pass 32s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326718
Ruff lint pass 8s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326682
Syntax sweep + import smoke (no GPU stack) pass 21s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326815
Full test suite (offline, mocked) pass 2m14s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326676
Golden fixtures cross-OS (macos-latest) pass 16s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326747
Golden fixtures cross-OS (ubuntu-latest) pass 23s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326771
Golden fixtures cross-OS (windows-latest) pass 1m5s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326779
Mypy (lenient baseline) pass 32s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326718
Ruff lint pass 8s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326682
Syntax sweep + import smoke (no GPU stack) pass 21s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27908114689/job/82580326815>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

169. assistant (2026-06-21T14:58:35.784Z)

All 7 CI checks green (full suite 2m14s, cross-OS golden fixtures, mypy, ruff). Per your standing authorization to merge verified-green good-engineering PRs, merging P0a (squash, matching the repo convention).

170. user (2026-06-21T14:58:42.549Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
=== master now ===
2afbadf..82a28f4 master -> origin/master
82a28f4 feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for config/.env/DB/output (default-OFF) (#293)
2afbadf fix(dryrun): install the watermark font in bootstrap + track the UI-trigger asks (#292)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

171. assistant (2026-06-21T14:58:51.715Z)

P0a is **merged** ? `origin/master` now has `82a28f4 ... P0a ...` (#293). The error was only `gh` trying to check out `master` locally for cleanup (it's held by another worktree); the merge itself succeeded. Let me finish cleanup (remote branch + worktree) and confirm.

172. user (2026-06-21T14:58:57.225Z)

```
=== PR #293 state ===
MERGED @ 2026-06-21T14:58:39Z
=== delete remote branch if present ===
error: unable to delete 'feat/packaging-p0a-app-home': remote ref does not exist
error: failed to push some refs to 'https://github.com/Khelsutra/badminton-highlight-indexer.git'
=== remove worktree ===
worktree removed
=== worktrees now ===
(no p0a worktree)
```

173. assistant (2026-06-21T14:59:24.385Z)

P0a is **fully shipped**: PR #293 **MERGED** to master (`82a28f4`), remote branch deleted, worktree cleaned up. Let me update the plan tracker + memory checkpoint, then continue to P0b.

174. user (2026-06-21T14:59:27.048Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

175. user (2026-06-21T14:59:44.290Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

176. assistant (2026-06-21T15:00:24.193Z)

P0a is a clean checkpoint (merged, reviewed, proof-validated). Continuing to **P0b** ? redact `GET /api/config` (it returns the entire `AppConfig` unauthenticated). I'll scope the Host-header/DNS-rebinding guard into P0c with the bind/CORS work (they're the same "network exposure" concern), keeping P0b tight: redaction + a serialization-safety test. Fresh worktree off the new

master, then read the surface.

177. user (2026-06-21T15:00:29.023Z)

```
branch 'feat/packaging-p0b-config-redact' set up to track 'origin/master'.
HEAD is now at 82a28f4 feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for
config/.env/DB/output (default-OFF) (#293)
=== branch/head ===
82a28f4 feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for config/.env/DB/output
(default-OFF) (#293)
=== /api/config consumers (engine + khelsutra SPA) ===
C:/Users/avidu/Projects/bhi-wt-p0b/backend/main.py:159:@app.get("/api/config")
=== AppConfig secret-shaped fields? ===
194: # (stream copy cuts at the nearest keyframe). 0 = legacy stream-copy at source res.
199: # emit their local candidate windows AS the rallies ? no provider, no API key, nothing
302: # application's audience tag. NO secrets here ? these are public identifiers.
305: cf_access_aud: str = "" # the CF Access application AUD tag the token must carry
307: # backend/main.py so importing the app does no filesystem I/O; default [""].
Credentials stay
316: **secrets are never stored here**, only endpoints / regions / file paths / source-
selectors
317: (boto3/google auth chains supply credentials). See ``backend/storage/``.
332: # One folder per video keyed by the MD5 video_id:
<root>/[<workspace_prefix>]/<video_id>/...
410: # Allow extra keys during transition so existing config.json files do not break.
422: def get(self, key: str, default: Any = None) -> Any:
```

178. assistant (2026-06-21T15:01:14.838Z)

`/api/config` is consumed only at its definition (the SPA uses `/api/health`), and AppConfig explicitly stores **no secrets today** (`cf\_access\_aud` is a public identifier) ? so this is **defensive/future-proofing**: redact any secret-shaped field and lock it with a test, without over-masking legit fields like `edge\_auth\_enabled`. Let me see how existing API tests set `global\_config` + drive the endpoint.

179. user (2026-06-21T15:01:16.615Z)

```
=== how tests set global_config + TestClient ===
1: """TestClient coverage for the FastAPI seam (previously zero ? the compiler SyntaxError
2: that broke `import backend.main` shipped green because nothing imported it)."""
10: pytest.importorskip("httpx") # TestClient needs httpx; skip cleanly where absent
11: from fastapi.testclient import TestClient
13: import backend.main as m
28: def env(tmp_path, monkeypatch):
31:     monkeypatch.setattr(m, "db_instance", db)
32:     monkeypatch.setattr(m, "global_config", cfg)
33:     monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
34:     client = TestClient(m.app)
35:     return client, db, tmp_path, monkeypatch
63:     r = client.get("/api/config")
108: def test_process_uses_configured_output_dir_not_literal_output(env, tmp_path, monkeypatch):
123:     mp.setattr(m, "global_config", cfg)
146: def test_compiler_and_reingest_safety_honor_output_dir_override(env, tmp_path, monkeypatch):
158:     mp.setattr(m, "global_config", cfg)
262:     r = client.get("/api/health")
276:     m.global_config.stitching.min_shot_count = 8
303:     m.global_config.stitching.min_shot_count = 8
318:     cfg = m.global_config
349: def test_parity_job_worker_vs_direct_cli_path(env, tmp_path, monkeypatch):
517:     mp.setattr(m.global_config.security, "allowed_video_root", str(root))
530: def test_cli_nonlocal_input_fetches_and_keeps_uri(tmp_path, monkeypatch):
536:     monkeypatch.setattr(m, "load_app_config", lambda *a, **k: cfg)
537:     monkeypatch.setattr(m, "Database", lambda *a, **k: db)
538:     monkeypatch.setattr(m, "global_config", cfg)
```

```
539: monkeypatch.setattr(m, "db_instance", db)
540: monkeypatch.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
541: monkeypatch.setattr(m.segmenter_registry, "get", lambda name: _CaptureSeg)
548: monkeypatch.setattr(m.storage, "fetch", fake_fetch)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

180. assistant (2026-06-21T15:01:22.239Z)

There's already an `/api/config` test at line 63. Let me read the fixture and that test to extend it.

181. user (2026-06-21T15:01:22.890Z)

```
20         """submit() runs the fn synchronously ? async jobs complete before the response
21         is inspected (full lifecycle coverage lives in test_async_jobs.py)."""
22
23         def submit(self, fn, *args, **kwargs):
24             fn(*args, **kwargs)
25
26
27     @pytest.fixture
28     def env(tmp_path, monkeypatch):
29         db = Database(str(tmp_path / "t.db"))
30         cfg = AppConfig()
31         monkeypatch.setattr(m, "db_instance", db)
32         monkeypatch.setattr(m, "global_config", cfg)
33         monkeypatch.setattr(m, "_get_executor", lambda: _InlineExecutor())
34         client = TestClient(m.app)
35         return client, db, tmp_path, monkeypatch
36
37
38     def _video(tmp_path, name="m.mp4"):
39         p = tmp_path / name
40         p.write_bytes(b"not-a-real-video-but-hashable")
41         return p
42
43
44     def _pass():
45         return [ValidationResult(validator_name="x", passed=True, message="ok", details={})]
46
47
48     def _fail():
49         return [ValidationResult(validator_name="blur", passed=False, message="blurry",
50 details={})]
51
52     class _FakeSeg:
53         def __init__(self, db, cfg):
54             self.db = db
55
56         def process_video(self, video_id, video_path, *a, **k):
57             sid = self.db.add_play_segment(video_id, 0.0, 5.0)
58             return [{"id": sid, "start_time": 0.0, "end_time": 5.0, "duration": 5.0,
59 "metadata": {}}]
60
61     def test_api_config_ok(env):
62         client = env[0]
63         r = client.get("/api/config")
64         assert r.status_code == 200
65         assert r.json()["sport"] == "badminton"
66
67
68     def test_process_success_with_mocked_segmenter(env):
69         """#16: submit returns 202 + job_id; the (inline) job carries the old response."""
```

```

70     client, db, tmp_path, mp = env
71     vid_file = _video(tmp_path)
72     mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_pass(), {}))
73     mp.setattr(m.segmenter_registry, "get", lambda name: _FakeSeg)
74     r = client.post("/api/process", json={"video_path": str(vid_file), "segmenter":
"x"})
75     assert r.status_code == 202
76     job = client.get(f"/api/jobs/{r.json()['job_id']}").json()
77     assert job["status"] == "done"
78     assert job["result"]["status"] == "success"
79     assert len(job["result"]["play_segments"]) == 1
80
81
82     def test_process_validation_failure_preserves_prior_segments(env):
83         client, db, tmp_path, mp = env
84         vid_file = _video(tmp_path)
85         vid = m.compute_video_id(str(vid_file))
86         db.add_video(vid, str(vid_file), "processed", [])
87         db.add_play_segment(vid, 1.0, 2.0) # prior good result
88         mp.setattr(m.registry, "run_all_with_context", lambda *a, **k: (_fail(), {}))
89         r = client.post("/api/process", json={"video_path": str(vid_file)})
90         assert r.status_code == 202
91         job = client.get(f"/api/jobs/{r.json()['job_id']}").json()

```

182. assistant (2026-06-21T15:01:48.802Z)

Clean picture: the existing `test\_api\_config\_ok` only checks `sport`, so redaction won't break it (sport isn't secret), and the SPA doesn't call `/api/config`. Implementing P0b ? a focused, well-bounded secret-redaction helper + tests.

183. assistant (2026-06-21T15:01:56.993Z)

"""Redact secret-shaped values before they leave the process (PACKAGING_PLAN.md P0b).

`GET /api/config` returns the whole `AppConfig` unauthenticated. Today AppConfig stores NO secrets by design (the Gemini key lives in the environment; `cf\_access\_aud` is a public identifier), so this is DEFENSIVE: if a credential-shaped field is ever added to the config, this masks it in the response and the accompanying test fails loudly rather than leaking it to any localhost process/browser tab.

The match list is deliberately NARROW so it never masks legitimate non-secret fields such as `edge\_auth\_enabled` (a boolean), `cf\_access\_aud` (a public id), or an ML `max\_tokens` knob.

```

from __future__ import annotations

import re
from typing import Any

#: Field-name patterns whose VALUE is masked. Narrow on purpose (see module docstring).
SECRET_KEY_RE = re.compile(
    r"(password|passwd|secret|credential|api[_-]?key|access[_-]?token|auth[_-]?token|refresh[_-]?token)",
    re.IGNORECASE,
)

REDACTED = "****REDACTED****"

```

```

def redact_secrets(obj: Any) -> Any:
    """Return a deep copy of `obj` with any value under a secret-shaped KEY replaced by
    :data:`REDACTED`. Recurses into dicts and lists; scalars and non-secret keys pass through.

```

A secret-named key masks its **WHOLE** value (even a nested dict/list ? e.g. a ``credentials`` block), so a structured secret can't slip through via recursion.

```
"""
if isinstance(obj, dict):
    out: dict[Any, Any] = {}
    for key, value in obj.items():
        if isinstance(key, str) and SECRET_KEY_RE.search(key):
            out[key] = REDACTED
        else:
            out[key] = redact_secrets(value)
    return out
if isinstance(obj, (list, tuple)):
    return [redact_secrets(item) for item in obj]
return obj
```

184. user (2026-06-21T15:01:57.013Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0b\backend\utils\redact.py (file state is current in your context ? no need to Read it back)

185. assistant (2026-06-21T15:02:00.050Z)

Now wire it into the ``/api/config`` handler.

186. user (2026-06-21T15:02:00.894Z)

```
155         db_instance.prune_segments(video_id, play_after=play_wm, cand_after=cand_wm)
156
157
158     # --- API Endpoints ---
159     @app.get("/api/config")
160     def get_config():
161         return global_config
162
```

187. user (2026-06-21T15:02:07.411Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

188. assistant (2026-06-21T15:02:13.578Z)

Add the import alongside the other ``backend.utils`` imports.

189. user (2026-06-21T15:02:14.717Z)

```
14:from backend.utils.app_home import env_file_to_load, resolve_output_dir
36:from backend.utils.hashing import compute_video_id # canonical file-identity hashing
37:from backend.utils.paths import resolve_under_root, safe_output_filename, validate_input_path
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

190. user (2026-06-21T15:02:19.589Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

191. assistant (2026-06-21T15:02:23.097Z)

Now the tests ? a unit test for the redactor and an endpoint test proving a secret field is masked while legit fields aren't.

192. assistant (2026-06-21T15:02:32.530Z)

```
"""Tests for the secret-redaction helper (PACKAGING_PLAN.md P0b)."""
```

```
from backend.utils.redact import REDACTED, redact_secrets
```

```
def test_masks_common_secret_field_names():
```

```
    src = {
        "gemini_api_key": "AIZA-xxx",
        "api_key": "sk-xxx",
        "access_token": "tok",
        "auth_token": "tok2",
        "refresh_token": "tok3",
        "client_secret": "shh",
        "password": "hunter2",
        "passwd": "hunter2",
        "db_credential": "pg://u:p@h",
    }
    out = redact_secrets(src)
    assert all(v == REDACTED for v in out.values()), out
```

```
def test_does_not_mask_legitimate_non_secret_fields():
```

```
    # The narrow matcher must NOT mask these real config fields / knobs.
    src = {
        "edge_auth_enabled": True,
        "cf_access_aud": "my-app-aud",
        "max_tokens": 1024,
        "sport": "badminton",
        "output_dir": "output",
        "tokenizer_name": "bert", # 'token' substring but not an auth token field
    }
    assert redact_secrets(src) == src
```

```
def test_recurse_into_nested_dicts_and_lists():
```

```
    src = {"storage": {"gcs": {"api_key": "k", "bucket": "b"}}, "items": [{"secret": "s", "n": 1}]}
    out = redact_secrets(src)
    assert out["storage"]["gcs"]["api_key"] == REDACTED
    assert out["storage"]["gcs"]["bucket"] == "b"
    assert out["items"][0]["secret"] == REDACTED
    assert out["items"][0]["n"] == 1
```

```
def test_secret_named_key_masks_whole_nested_value():
```

```
    # A structured secret under a secret-named key is masked wholesale (no recursion leak).
    src = {"credentials": {"user": "u", "pass": "p"}}
    assert redact_secrets(src) == {"credentials": REDACTED}
```

```
def test_scalars_and_non_dict_inputs_pass_through():
```

```
    assert redact_secrets(42) == 42
    assert redact_secrets("plain") == "plain"
    assert redact_secrets([1, 2, 3]) == [1, 2, 3]
```

```
def test_does_not_mutate_input():
```

```
    src = {"api_key": "k", "nested": {"token_access": "t"}}
    snapshot = {"api_key": "k", "nested": {"token_access": "t"}}
    redact_secrets(src)
    assert src == snapshot
```

193. user (2026-06-21T15:02:32.564Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0b\tests\test_redact.py (file

state is current in your context ? no need to Read it back)

194. assistant (2026-06-21T15:03:04.233Z)

Add the endpoint redaction test (proves the wiring: a secret-shaped field is masked, non-secret fields still served).

195. user (2026-06-21T15:03:08.922Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\tests\test_api_endpoints.py has been updated successfully. (file state is current in your context ? no need to Read it back)

196. user (2026-06-21T15:03:30.104Z)

```
..... [100%]
37 passed in 12.11s
=== RUFF+MYPY ===
All checks passed!
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

197. assistant (2026-06-21T15:03:34.159Z)

P0b tests green, ruff+mypy clean. Full suite + coverage gate.

198. user (2026-06-21T15:05:36.010Z)

```
backend\main.py          424    127    70%
backend\utils\redact.py    16     0   100%
TOTAL                   9650   1436    85%
Required test coverage of 70% reached. Total coverage: 85.12%
1 failed, 1238 passed, 7 skipped, 7 warnings in 117.24s (0:01:57)
RESULT: FAIL - pytest exit 1
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

199. assistant (2026-06-21T15:05:41.986Z)

One failure (coverage fine). Let me find it ? likely a test asserting `/api/config`'s exact shape.

200. user (2026-06-21T15:07:12.667Z)

```
E      AssertionError: assert False
tests\test_config.py:328: AssertionError
FAILED tests/test_config.py::test_config_roundtrip_and_typed_reaches_pipeline
1 failed, 1238 passed, 7 skipped in 86.42s (0:01:26)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

201. user (2026-06-21T15:07:15.926Z)

```
300         assert cfg.get("ai_provider", "ignored-default") == "gemini"
301         assert cfg.get("does_not_exist", "fallback") == "fallback"
302
303         # leaf keys inside any section resolve via the dynamic section scan (not a hardcoded
304         # list), so a future top-level section's keys would be discoverable too
305         assert cfg.get("min_resolution_width") == 1280 # lives inside `validation`
306         assert cfg.get("begin_padding") == 0.5        # lives inside `stitching`
307
308         # legacy chained access still works
309         assert cfg.get("indexing", {}).get("default_segmenter") == "yolo_hybrid"
310
311
312     def test_config_roundtrip_and_typed_reaches_pipeline(monkeypatch):
```

```

313         """Positive for Item 2: save_default_config + load_config round-trips;
314         typed AppConfig is what reaches pipeline code (e.g. main uses global_config as
AppConfig).
315         """
316         with tempfile.TemporaryDirectory() as td:
317             p = os.path.join(td, "roundtrip.json")
318             save_default_config(p)
319             cfg = load_config(p)
320             assert isinstance(cfg, AppConfig)
321             assert cfg.sport == "badminton"
322
323         # positive that typed is used in pipeline code (global_config is set at startup in
main());
324         # monkeypatch here so the test does not depend on server initialization)
325         import backend.main as main_module
326         from backend.main import get_config
327         monkeypatch.setattr(main_module, "global_config", cfg)
328         assert isinstance(get_config(), AppConfig)
329

```

202. assistant (2026-06-21T15:07:32.746Z)

The redaction made `get\_config()` return a dict (the API response), so a test asserting the handler returns the raw `AppConfig` model is now outdated. I'll update it to assert the new contract ? the handler returns a redacted dict while `global\_config` stays the typed model.

203. user (2026-06-21T15:07:36.954Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\tests\test_config.py has been updated successfully. (file state is current in your context ? no need to Read it back)

204. user (2026-06-21T15:09:16.773Z)

```

..... [100%]
21 passed in 6.03s
=== full suite ===
1239 passed, 7 skipped in 83.78s (0:01:23)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

205. assistant (2026-06-21T15:09:40.886Z)

Full suite green (1239 passed). P0b is small and purely subtractive (redaction), but to honor the per-PR review norm I'll run one focused adversarial reviewer on the diff before committing.

206. user (2026-06-21T15:10:58.949Z)

I have everything needed. Here are my findings.

Findings

****[minor] backend/utils/redact.py:20 ? Under-masking: several realistic secret-shaped names slip the regex.**** The narrow allowlist misses `private\_key`, `bearer`/`bearer\_token`, `access\_key`, `aws\_access\_key\_id`, `signing\_key`/`encryption\_key`, `session\_token`, `gh\_token`, `authorization`, `cookie`. Notably `api\_token` and `gemini\_key` (a plausible variant of the very example the docstring/test centers on) both MISS ? only `api\_key`/`api-key`/`access\_token` match, not `\*\_token` generally or `\*\_key` generally. This is a defensive future-proofing layer, so a miss only bites if someone later adds such a field AND it isn't caught by the accompanying test. Acceptable for P0b's stated narrow-on-purpose intent, but worth widening. Concrete fix: add alternations `|token|private[\_]?key|bearer|access[\_]?key|signing[\_]?key|encryption[\_]?key|authorization|cookie` (accepting that bare `token` would also mask `auth\_token` variants and could catch a future non-secret `\*token` field ? tune against the real field list).

****[nit] backend/utils/redact.py:42-43 ? tuples are silently converted to lists.****
`redact\_secrets((1,2))` returns `[1,2]`. Irrelevant for the `/api/config` JSON path (Pydantic `model\_dump(mode="json")` already emits lists), so no live impact; only matters if the helper is reused elsewhere on tuple-bearing data. Concrete fix: none required for this PR; note in the docstring that tuples normalize to lists, or `return type(obj)(...)` if tuple-preservation is ever wanted.

****No over-masking (verified):**** the regex matches ****zero**** of the real AppConfig field names ? enumerated all fields across AppConfig + ValidationConfig/relevance, TrackNet, WasbIndex, Fusion, Indexing, Output, Watermark, Stitching, Logging, Security, Storage, Deployment, Billing. In particular `cf\_access\_aud` (public id), `edge\_auth\_enabled` (bool), `auth`-adjacent names, `access\_\*` (e.g. none present), `\*token` (none present), `credential`/`secret` substrings (none present), and `\*\_key` knobs (none present) all pass through unmasked. The `cf\_access\_aud` field contains "access" but not "access_token", so it correctly does NOT match ? no false positive.

****Serialization equivalence (verified):**** `global\_config.model\_dump(mode="json")` is byte-equal to the old `return global\_config` path (`jsonable\_encoder(model)`) for every non-secret field ? confirmed by direct comparison. All config fields are str/int/float/bool/list/dict/`Literal`; `Literal` values serialize as their plain string in both paths, and there are no `Path`/`datetime`/`Enum`/`set` fields that would diverge. So the response shape for real fields is unchanged.

****Test validity (verified):**** `AppConfig(gemini\_api\_key=...)` is accepted via `model\_config extra="allow"` (models.py:411), and `model\_dump(mode="json")` surfaces the extra key ? so the injected `gemini\_api\_key` is genuinely present pre-redaction and the test's `==`
`\*\*\*REDACTED\*\*\*` assertion exercises a real masking, not a no-op.

****Leak-path analysis (clean):**** A secret-named key masks its whole value before recursing (redact.py:37-38), so a structured secret (nested dict/list under e.g. `credentials`) can't escape via recursion ? covered by `test\_secret\_named\_key\_masks\_whole\_nested\_value`. Lists of dicts recurse correctly. The `None`-guard returns `{}` (main.py:166-167), leaking nothing. The one genuine residual gap is values where the secret is in the VALUE not the KEY ? e.g. a connection string in a non-secret-named field like `storage.gcs = {"endpoint\_url": "https://user:pass@host"}` or `input\_root="gs://...@..."`. Redaction is key-name-based only, so such a value would pass through. This matches the documented design (config "stores no secrets? boto3/google auth chains supply credentials", models.py:317) and is out of scope for a key-name redactor; flagging only for completeness, not as a defect of this PR.

****Verdict: LGTM with minor follow-up**** ? correct, no over-masking, serialization byte-compatible, test valid; the only substantive note is the deliberately narrow regex misses some plausible future secret names (`private\_key`, `bearer`, `\*\_token`, `access\_key`), which is consistent with the PR's stated narrow-on-purpose stance but worth widening before a real credential field lands.

(Note: I did not run `git pull`/`run\_all\_tests.py`/ruff/mypy here ? this was a scoped read-only adversarial review of the four named files; full pre-LGTM suite verification per your gate is still owed before sign-off.)
agentId: a83eb777e68482055 (use SendMessage with to: 'a83eb777e68482055' to continue this agent)
<usage>subagent_tokens: 58496
tool_uses: 8
duration_ms: 67206</usage>

207. assistant (2026-06-21T15:11:28.066Z)

Solid review ? no over-masking, serialization-compatible, test valid. The one substantive note (regex too narrow, misses `private\_key`/`bearer`/`\*\_token`/`gemini\_key`) is worth acting on since masking-more is the safe direction for an unauthenticated endpoint. Let me first verify no real field ends in `\_key`/`\_token` (so widening can't over-mask), then widen + add a full-config regression guard.

208. user (2026-06-21T15:11:30.587Z)

=== any real field ending _key / _token / containing bearer/private_key/access_key? ===
308: # OFF ? Cf-Access is a header, not a cookie.)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

209. assistant (2026-06-21T15:11:35.401Z)

Confirmed ? no real field would be over-masked (the only "cookie" hit is a comment). Widening the matcher and adding a full-`AppConfig` regression guard against over-masking.

210. user (2026-06-21T15:11:41.337Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\backend\utils\redact.py has been updated successfully. (file state is current in your context ? no need to Read it back)

211. assistant (2026-06-21T15:11:51.970Z)

Now strengthen the tests: add the new secret names and a full-`AppConfig` over-masking guard.

212. user (2026-06-21T15:11:57.672Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\tests\test_redact.py has been updated successfully. (file state is current in your context ? no need to Read it back)

213. user (2026-06-21T15:12:06.345Z)

The file C:\Users\avidu\Projects\bhi-wt-p0b\tests\test_redact.py has been updated successfully. (file state is current in your context ? no need to Read it back)

214. user (2026-06-21T15:12:26.141Z)

```
..... [100%]
59 passed in 11.89s
=== ruff+mypy ===
All checks passed!
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

215. assistant (2026-06-21T15:12:30.599Z)

All green including the full-`AppConfig` over-masking guard. Final full-suite run, then commit/PR.

216. user (2026-06-21T15:14:29.386Z)

```
backend\utils\redact.py          16      0   100%
TOTAL                          9650  1436   85%
Required test coverage of 70% reached. Total coverage: 85.12%
1240 passed, 7 skipped, 7 warnings in 113.40s (0:01:53)
  RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

217. assistant (2026-06-21T15:14:34.599Z)

Full suite green (1240 passed), coverage 85.12%, redact.py 100%. Committing P0b (explicit paths) and opening the PR.

218. user (2026-06-21T15:14:51.209Z)

```
=== status (must be only P0b files) ===
M backend/main.py
M tests/test_api_endpoints.py
M tests/test_config.py
?? backend/utils/redact.py
?? tests/test_redact.py
warning: in the working copy of 'backend/utils/redact.py', LF will be replaced by CRLF the next
time Git touches it
warning: in the working copy of 'tests/test_redact.py', LF will be replaced by CRLF the next
```

```
time Git touches it
=== HEAD ===
07157cd feat(pkg): P0b ? redact secret-shaped fields from the unauthenticated GET /api/config
=== push ===
branch 'feat/packaging-p0b-config-redact' set up to track 'origin/feat/packaging-p0b-config-redact'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
 * [new branch]      feat/packaging-p0b-config-redact -> feat/packaging-p0b-config-redact
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

219. user (2026-06-21T15:15:06.748Z)

```
https://github.com/Khelsutra/badminton-highlight-indexer/pull/295
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

220. assistant (2026-06-21T15:15:10.694Z)

```
**PR #295 opened: [feat(pkg): P0b ? redact
`/api/config`](https://github.com/Khelsutra/badminton-highlight-indexer/pull/295).** Watching CI
to green.
```

221. user (2026-06-21T15:17:57.259Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    22s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27908601199/job/82581659551
Full test suite (offline, mocked)               pass    2m29s  https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27908601199/job/82581659548
Golden fixtures cross-OS (macos-latest)          pass    17s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27908601199/job/82581659603
Golden fixtures cross-OS (ubuntu-latest)         pass    24s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27908601199/job/82581659638
Golden fixtures cross-OS (windows-latest)        pass    49s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27908601199/job/82581659585
Mypy (lenient baseline)                          pass    36s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27908601199/job/82581659558
Ruff lint                                         pass    10s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27908601199/job/82581659610
Syntax sweep + import smoke (no GPU stack)      pass    22s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27908601199/job/82581659551
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

222. assistant (2026-06-21T15:18:00.296Z)

All 7 CI checks green. Merging P0b and cleaning up.

223. user (2026-06-21T15:18:08.990Z)

```
failed to run git: fatal: 'master' is already used by worktree at
'C:\Users\avidu\Projects\badminton-highlight-indexer/.claude/worktrees/tier1-windowing'
```

```
0c0e064..fc78fce master -> origin/master
=== master ===
fc78fce feat(pkg): P0b ? redact secret-shaped fields from the unauthenticated GET /api/config
(#295)
p0b worktree removed
=== PR #295 ===
MERGED
```

224. assistant (2026-06-21T15:18:30.725Z)

P0b ****merged**** (master `fc78fce`). Updating the tracker + memory, then continuing to P0c.

225. user (2026-06-21T15:18:35.645Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

226. user (2026-06-21T15:18:48.375Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

227. assistant (2026-06-21T15:20:03.011Z)

Now ****P0c ? config/profile-driven bind + DNS-rebinding (Host-header) guard****. This is the meatiest P0 item (network exposure), so I'll design it carefully to stay behaviour-neutral (default `127.0.0.1`, TestClient's `testserver` host preserved). Fresh worktree, then read the CORS/bind/SecurityConfig surface.

228. user (2026-06-21T15:20:07.697Z)

```
branch 'feat/packaging-p0c-bind-host' set up to track 'origin/master'.
HEAD is now at fc78fce feat(pkg): P0b ? redact secret-shaped fields from the unauthenticated GET
/api/config (#295)
=== head ===
fc78fce feat(pkg): P0b ? redact secret-shaped fields from the unauthenticated GET /api/config
(#295)
=== CORS + bind + TrustedHost ===
26:from fastapi.middleware.cors import CORSMiddleware
55:def _cors_origins_from_env(env: Optional[Dict[str, str]] = None) -> List[str]:
56:    """M9: CORS allow-origins from the `RALLY_CORS_ALLOW_ORIGINS` env var (comma-
separated),
60:    raw = (src.get("RALLY_CORS_ALLOW_ORIGINS", "") or "").strip()
69:app.add_middleware(
70:    CORSMiddleware,
71:    allow_origins=_cors_origins_from_env(),
608:    parser.add_argument("--port", type=int, default=8000, help="Port to run the FastAPI
server on")
782:    uvicorn.run(app, host="127.0.0.1", port=args.port)
=== SecurityConfig fields ===
280:class SecurityConfig(BaseModel):
281:    # When set, /api/process video_path must resolve under this directory (hosted/tenant
283:    allowed_video_root: Optional[str] = None
287:    # arrive with PR-3 identity scoping. ? In hosted mode, set allowed_video_root to
297:    # DEFAULT-OFF: edge_auth_enabled=False ? no JWT is required/verified (the local single-
user
303:    edge_auth_enabled: bool = False
305:    cf_access_aud: str = ""          # the CF Access application AUD tag the token must carry
```

229. user (2026-06-21T15:20:12.266Z)

```
55     def _cors_origins_from_env(env: Optional[Dict[str, str]] = None) -> List[str]:
56         """M9: CORS allow-origins from the `RALLY_CORS_ALLOW_ORIGINS` env var (comma-
separated),
57         default `[*]`. Read from the ENV (not a cwd `config.json`) so importing
`backend.main`
58         does zero filesystem I/O ? a hosted/multi-tenant deploy sets this to lock CORS to
its origin."""
59         src = os.environ if env is None else env
60         raw = (src.get("RALLY_CORS_ALLOW_ORIGINS", "") or "").strip()
61         origins = [o.strip() for o in raw.split(",") if o.strip()]
62         return origins or ["*"]
63
```

```

64
65     # CORS for the web UI. allow_origins='' with allow_credentials=True is rejected by
browsers per the
66     # fetch spec + is unsafe to carry into multi-tenant; this tool uses no cookies, so
credentials stay
67     # off. The origin list is configurable via the env var above (default [""]) so a hosted
deploy can
68     # lock it to its origin; the middleware is fixed at startup.
69     app.add_middleware(
70         CORSMiddleware,
71         allow_origins=_cors_origins_from_env(),
72         allow_credentials=False,
73         allow_methods=["*"],
74         allow_headers=["*"],
75     )
76
77     # Serves static frontend files directly (now under api/ per approved structure).
78     # PACKAGING_PLAN.md P0a: resolve PACKAGE-relative (not CWD-relative) so launching from
an
79     # arbitrary directory finds the shipped assets instead of creating a stray
./backend/api/static.
80     # (The legacy in-engine dashboard; the real SPA gets bundled here in P1a.)

```

230. user (2026-06-21T15:20:16.464Z)

```

772         print(f"[REPORT] Machine JSON      : {paths.get('json')}")
773         return
774
775     # Start Local Server Mode
776     if args.server or not args.video_path:
777         print(f"[SERVER] Starting local FastAPI server on
http://localhost:{args.port}...")
778         # Resume self-scheduled work orphaned by the restart (EXECUTION_POLICY.md P3):
the
779         # in-process wake-timer for any parked 'waiting' job died with the old process,
so
780         # kick one drain ? it reclaims due jobs and re-arms the wake-timer for the rest.
781         _get_executor().submit(_drain_due_jobs)
782         uvicorn.run(app, host="127.0.0.1", port=args.port)
783

```

231. assistant (2026-06-21T15:20:58.174Z)

Clear. Implementing P0c: a `TrustedHostMiddleware` DNS-rebinding guard (env-driven allowlist defaulting to loopback + TestClient's `testserver`), and a profile-aware bind-host resolver (default `127.0.0.1`; `0.0.0.0` only under edge-auth). Adding the import + helpers + middleware.

232. user (2026-06-21T15:21:00.060Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

233. user (2026-06-21T15:21:18.078Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

234. assistant (2026-06-21T15:21:20.848Z)

Now the `--host` CLI arg and the bind-host wiring at server start.

235. user (2026-06-21T15:21:23.341Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\main.py has been updated successfully. (file

state is current in your context ? no need to Read it back)

236. user (2026-06-21T15:21:32.285Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

237. assistant (2026-06-21T15:21:43.362Z)

Now tests for the bind resolver, the host allowlist, and the middleware's actual reject/allow behavior via TestClient.

238. assistant (2026-06-21T15:21:53.914Z)

```
"""Bind-address resolution + DNS-rebinding Host-header guard (PACKAGING_PLAN.md P0c)."""
```

```
import pytest
```

```
pytest.importorskip("httpx") # TestClient needs httpx
from fastapi.testclient import TestClient
```

```
import backend.main as m
from backend.config import AppConfig
```

```
@pytest.fixture(autouse=True)
def _clear_env(monkeypatch):
    for var in ("RALLY_BIND_HOST", "RALLY_ALLOWED_HOSTS"):
        monkeypatch.delenv(var, raising=False)
```

---- bind host resolution ----

```
def test_bind_defaults_to_localhost():
    assert m._resolve_bind_host(AppConfig(), None) == "127.0.0.1"
```

```
def test_bind_cli_host_wins(monkeypatch):
    monkeypatch.setenv("RALLY_BIND_HOST", "10.0.0.5")
    cfg = AppConfig()
    cfg.security.edge_auth_enabled = True
    assert m._resolve_bind_host(cfg, "192.168.1.9") == "192.168.1.9" # --host beats env +
profile
```

```
def test_bind_env_beats_profile(monkeypatch):
    monkeypatch.setenv("RALLY_BIND_HOST", "0.0.0.0")
    assert m._resolve_bind_host(AppConfig(), None) == "0.0.0.0"
```

```
def test_bind_edge_auth_profile_widens_to_all_interfaces():
    cfg = AppConfig()
    cfg.security.edge_auth_enabled = True
    assert m._resolve_bind_host(cfg, None) == "0.0.0.0"
```

```
def test_bind_no_config_is_localhost():
    assert m._resolve_bind_host(None, None) == "127.0.0.1"
```

---- allowed-hosts resolution ----

```
def test_allowed_hosts_default_is_loopback_plus_testserver():
```

```

hosts = m._allowed_hosts_from_env({})
assert "127.0.0.1" in hosts and "localhost" in hosts and "testserver" in hosts
assert "*" not in hosts

def test_allowed_hosts_env_override():
    assert m._allowed_hosts_from_env({"RALLY_ALLOWED_HOSTS": "app.example.com, 1.2.3.4"}) == [
        "app.example.com", "1.2.3.4"]

def test_allowed_hosts_wildcard_for_edge():
    assert m._allowed_hosts_from_env({"RALLY_ALLOWED_HOSTS": "*"}) == ["*"]

```

---- the middleware actually rejects a foreign Host (DNS-rebinding) ----

```

def test_foreign_host_header_is_rejected():
    client = TestClient(m.app)
    r = client.get("/api/health", headers={"host": "evil.attacker.com"})
    assert r.status_code == 400 # TrustedHostMiddleware: Invalid host header

def test_loopback_host_header_is_accepted():
    client = TestClient(m.app)
    # default TestClient host is "testserver" (allowed); an explicit loopback host (with port)
    # too.
    assert client.get("/api/health").status_code == 200
    assert client.get("/api/health", headers={"host": "127.0.0.1:8000"}).status_code == 200

```

239. user (2026-06-21T15:21:53.930Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_bind_and_host.py
(file state is current in your context ? no need to Read it back)

240. user (2026-06-21T15:22:05.283Z)

```

..... [100%]
10 passed in 6.03s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

241. assistant (2026-06-21T15:22:11.028Z)

P0c tests green (incl. the DNS-rebinding 400). The critical check is whether the global middleware breaks any existing TestClient test ? running the full suite.

242. user (2026-06-21T15:23:42.995Z)

```

..... [ 95%]
..... [100%]
1270 passed, 7 skipped in 87.71s (0:01:27)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

243. assistant (2026-06-21T15:23:51.646Z)

Full suite green (1270 passed) ? the global Host guard broke nothing (TestClient's `testserver` is allowlisted). Now lint/type + coverage.

244. user (2026-06-21T15:26:20.899Z)

```

=== RUFF ===

```

```

All checks passed!
=== MYPY ===
backend\sports\__init__.py:7: note: By default the bodies of untyped functions are not checked,
consider using --check-untyped-defs [annotation-unchecked]
backend\pipeline\validators\base.py:36: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 1 source file
=== COVERAGE ===
backend\main.py                                445      130      71%
TOTAL                                           9693     1442     85%
Required test coverage of 70% reached. Total coverage: 85.12%
  RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

245. assistant (2026-06-21T15:26:40.670Z)

Ruff clean, mypy clean (the notes are informational on unrelated modules), coverage 85.12%. I spotted one footgun myself: an edge-auth/`0.0.0.0` deploy that forgets `RALLY\_ALLOWED\_HOSTS` would have the loopback-default allowlist reject all real (public-Host) requests. Let me make the bind-time warning call that out, then run a focused review.

246. user (2026-06-21T15:26:44.394Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

247. user (2026-06-21T15:30:23.397Z)

`PACKAGING\_PLAN.md` doesn't exist in the worktree (it's a planning doc referenced in the prompt but not in-repo), and `RALLY\_ALLOWED\_HOSTS` is documented only in the docstring + the bind-time print. There's no standalone operator doc. The footgun surfacing rests entirely on the runtime print warning.

I have everything I need. Let me do a final precise check on the `getattr` null-safety in `\_resolve\_bind\_host` (concern #4) ? confirm `\_resolve\_bind\_host(None, None)` works (already tested, returns `127.0.0.1`) and that `getattr(sec, "edge\_auth\_enabled", False)` when `sec` is `None` returns `False` safely.

`getattr(None, "edge\_auth\_enabled", False)` ? `False` (safe; the test `test\_bind\_no\_config\_is\_localhost` proves it). `getattr(config, "security", None)` when `config=None` ? `None`. Both chains are null-safe.

Let me compile the findings.

Review findings

[minor] backend/main.py:90 (`\_allowed\_hosts\_from\_env`) + Starlette trustedhost.py:40 ? IPv6 loopback Host headers are silently rejected. The default allowlist includes `"::1"` and `"::1]"`, but Starlette's `TrustedHostMiddleware` matches with `host = headers.get("host", "").split(":")[0]`. A real IPv6 Host header is `"::1]:8000"` ? `".split(":")[0]"` yields `"::1]"` (and bare `"::1]"` ? `""`), neither of which matches any allowlist entry. Empirically verified: `"::1]:8000"`, `"::1]"`, and `"::1]"` all return **400**. So if a user binds/visits the appliance over IPv6 loopback, every request is wrongly rejected. The `"::1" / `"::1]"` entries in the default are dead ? they can never match given Starlette's port-stripping. Fix: document that IPv6 loopback is unsupported by this guard, or (cleaner) front it with a host-normalizer. The prompt flagged this as minor and it is ? IPv4 loopback (the actual default bind `127.0.0.1`) works fine.

[minor] backend/config/startup.py:91-139 vs backend/main.py:96-100 ? the edge-auth exposure footgun is surfaced only as a soft print, not enforced like its sibling checks. When `edge\_auth\_enabled=True`, `\_resolve\_bind\_host` returns `0.0.0.0` (main.py:114-115) but the TrustedHost allowlist is frozen at import from `RALLY\_ALLOWED\_HOSTS` only (main.py:98) ? it

cannot see config. If `RALLY\_ALLOWED\_HOSTS` is unset, the loopback-only default 400s ****100%**** of forwarded (public-Host) requests from the edge gate. The `\_exposure\_checks` startup gate already FATALs on a missing `allowed\_video\_root` and missing `cf\_team\_domain`/`cf\_access\_aud`, but does ****not**** check the host allowlist ? so a half-configured exposure (host-allowlist unset) starts "healthy" then 400s every real request. The only mitigation is the bind-time `print()` warning (main.py:823-827), which does correctly fire for `0.0.0.0` (verified: `0.0.0.0` is not in the suppress-tuple). Fix: add an `include\_exposure` startup check that warns/errors when `edge\_auth\_enabled` and `RALLY\_ALLOWED\_HOSTS` is unset/loopback-only, mirroring the existing `EXPOSED\_` checks ? that's the cleaner option than the env-only docstring note. Not a blocker because the default path is unaffected and the print does surface it.

****[nit]** backend/main.py:822 ? default server-start banner string changed from `localhost` to the resolved bind host. ****** Was `http://localhost:{port}`; now `http://{bind\_host}:{port}` which prints `http://127.0.0.1:8000` on the default path. The ***bind*** is byte-identical (`uvicorn.run` got `127.0.0.1` before too), but the user-visible banner text changed `localhost`?`127.0.0.1`. Purely cosmetic; flagging only because the contract says "byte-identical." No functional impact.

Items checked and CLEARED (no issue)

- ****Concern 1 ? legitimate access:**** `testserver` is allowlisted (main.py:90) ? full TestClient suite passes (247 API tests + 10 new = green). Loopback Host ****with port**** `127.0.0.1:8000` ? ****200**** (Starlette strips the port at trustedhost.py:40). Health, static `/`, and a good-host CORS preflight all return 200. No legitimate IPv4 path is wrongly 400'd.
- ****Concern 3 ? middleware order & CORS:**** Verified via `build\_middleware\_stack` that the wrap order is ServerError ? ****TrustedHost ? CORS**** ? router, i.e. TrustedHost is ****outermost**** and runs ***before*** CORS (FastAPI prepends; the stack is built inner-first via `reversed`). Empirically, a preflight `OPTIONS` from a disallowed Host returns 400 with ****no `Access-Control-Allow-Origin` header leaked****. `www\_redirect=False` is correct ? it prevents a 307 to a `www.` host (an open-redirect-ish surprise) and just 400s instead.
- ****Concern 2 (precedence) ? `\_resolve\_bind\_host` order** cli > env > profile > 127.0.0.1:****** correct and safe; matches the four tests (`test\_bind\_cli\_host\_wins`, `test\_bind\_env\_beats\_profile`, `test\_bind\_edge\_auth\_profile\_widens\_to\_all\_interfaces`, `test\_bind\_defaults\_to\_localhost`).
- ****Concern 4 ? null-safety:**** `getattr(config, "security", None)` then `getattr(sec, "edge\_auth\_enabled", False)` is null-safe for `config=None` AND `security=None` ? returns `127.0.0.1` (proven by `test\_bind\_no\_config\_is\_localhost`). The `--host` default `None` flows through `if cli\_host and cli\_host.strip()` correctly.
- ****Hard contract (byte-identity of default path):**** Holds. Default bind = `127.0.0.1` (unchanged), default allowlist preserves `testserver`, import does zero filesystem I/O (env-only), and the full existing TestClient suite is green. The only default-path delta is the cosmetic banner string above.

****Verdict:**** Not a blocker ? the hard contract holds and the default path is behaviorally byte-identical; ship-able, but consider promoting the edge-auth host-allowlist footgun (finding 2) to a startup check for parity with the existing `EXPOSED\_` gates, and document the IPv6-loopback limitation (finding 1).

agentId: af2b29f0bbdf75b59 (use SendMessage with to: 'af2b29f0bbdf75b59' to continue this agent)
<usage>subagent_tokens: 75258
tool_uses: 18
duration_ms: 201816</usage>

248. assistant (2026-06-21T15:30:44.232Z)

Thorough review ? verdict not-a-blocker, default path byte-identical. Two minor items worth folding: (1) document the IPv6-loopback limitation (Starlette's host parsing can't match `:::1`), and (2) promote the edge-auth + missing-`RALLY\_ALLOWED\_HOSTS` footgun to a startup check for parity with the existing `EXPOSED\_` gates. Let me check the startup-checks structure to add #2 cleanly.

249. user (2026-06-21T15:30:46.453Z)

19:behind an edge gate (the owner flips ``security.edge\_auth\_enabled`` on for the Cloudflare-Access

22:rest of the posture at ****boot****: the ``allowed\_video\_root`` path-jail is set (else an authenticated


```

30:a strict **no-op** ? exactly the pre-M5 behaviour; with ``edge_auth_enabled=False`` (today's
default)
42:LEVEL_WARN = "warn"
54:class StartupCheck:
62:class StartupCheckError(RuntimeError):
63:    """Raised by ``enforce_startup_checks`` when an active profile fails a FATAL check."""
65:    def __init__(self, failures: list[StartupCheck]):
75:        only downgrades to a WARNING, never a hard failure (the storage init is the
authority)."""
82:        (whether PyJWT is installed in the test env must not change the WARN's presence)."""
90:    def _exposure_checks(config: Any) -> list[StartupCheck]:
91:        """Exposure-safety posture for an *exposed* HTTP server. Gated on
``security.edge_auth_enabled``
97:        if not sec or not getattr(sec, "edge_auth_enabled", False):
100:            out: list[StartupCheck] = []
103:            # a tenant can point the engine at an ARBITRARY local file (allowed_video_root unset =
"any local
105:            root = getattr(sec, "allowed_video_root", None)
107:            out.append(StartupCheck(
108:                LEVEL_ERROR, "EXPOSED_NO_VIDEO_ROOT",
109:                "security.edge_auth_enabled is ON (engine exposed) but
security.allowed_video_root is "
111:                "Set allowed_video_root to a jail directory before exposing the engine.",
117:                team = (getattr(sec, "cf_team_domain", "") or "").strip()
118:                aud = (getattr(sec, "cf_access_aud", "") or "").strip()
120:                missing = " + ".join(n for n, v in (("cf_team_domain", team), ("cf_access_aud",
aud)) if not v)
121:                out.append(StartupCheck(
122:                    LEVEL_ERROR, "EXPOSED_EDGE_AUTH_INCOMPLETE",
123:                    f"security.edge_auth_enabled is ON but {missing} is not configured ? the Cf-
Access JWT "
124:                    "cannot be verified (every request would fail). Set security.cf_team_domain "
125:                    "(e.g. https://<team>.cloudflareaccess.com) and security.cf_access_aud (the CF
Access "
130:                    # missing, every request 500s at verify time ? a boot WARN beats a per-request 500.
Warning, not
133:                    out.append(StartupCheck(
134:                        LEVEL_WARN, "EXPOSED_AUTH_EXTRA_MISSING",
135:                        "security.edge_auth_enabled is ON but PyJWT[crypto] is not importable ? Cf-
Access "
148:) -> list[StartupCheck]:
155:    DEFAULT-OFF: no profile ? no profile checks; ``edge_auth_enabled=False`` ? no exposure
checks.
159:    checks: list[StartupCheck] = []
162:    # runs with no profile set), gated on edge_auth_enabled (default-OFF). Server-only.
179:        checks.append(StartupCheck(
180:            LEVEL_WARN, "STORAGE_UNUSUAL",
188:            checks.append(StartupCheck(
189:                LEVEL_WARN, "CUDA_NOT_VISIBLE",
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

250. assistant (2026-06-21T15:30:52.620Z)

The `\_exposure\_checks` structure is clean ? I'll add an `EXPOSED\_HOST\_ALLOWLIST\_LOOPBACK` WARN mirroring the existing `EXPOSED\_\*` checks. Let me read the exact block.

251. user (2026-06-21T15:30:53.464Z)

```

90    def _exposure_checks(config: Any) -> list[StartupCheck]:
91        """Exposure-safety posture for an *exposed* HTTP server. Gated on
``security.edge_auth_enabled``
92        (default False ? ``[]`` = strict no-op, today's behaviour). When edge auth is ON ?
the "I am
93        exposing this engine behind the CF-Access gate" signal ? the gate alone isn't
enough: verify the

```

```

94         rest of the posture at boot so a half-configured exposure fails fast instead of
leaking / 500-ing
95         per request. Config-only (no IO beyond the ``find_spec`` auth-extra probe)."""
96         sec = getattr(config, "security", None)
97         if not sec or not getattr(sec, "edge_auth_enabled", False):
98             return [] # DEFAULT-OFF: edge auth off ? no posture enforcement (byte-identical
to today).
99
100         out: list[StartupCheck] = []
101
102         # (1) The path-jail is load-bearing once authenticated tenants can hit /api/process:
WITHOUT it,
103         # a tenant can point the engine at an ARBITRARY local file (allowed_video_root unset
= "any local
104         # file", the single-user default). Fatal for an exposed engine.
105         root = getattr(sec, "allowed_video_root", None)
106         if not root or not str(root).strip():
107             out.append(StartupCheck(
108                 LEVEL_ERROR, "EXPOSED_NO_VIDEO_ROOT",
109                 "security.edge_auth_enabled is ON (engine exposed) but
security.allowed_video_root is "
110                 "unset ? an authenticated tenant could make /api/process read an arbitrary
local file. "
111                 "Set allowed_video_root to a jail directory before exposing the engine.",
112             ))
113
114         # (2) Edge auth cannot verify a token without the team domain + AUD. Today that only
fails on the
115         # FIRST request (auth.resolve_tenant); lift it to BOOT so a misconfigured gate never
starts
116         # "healthy" then 401s every request.
117         team = (getattr(sec, "cf_team_domain", "") or "").strip()
118         aud = (getattr(sec, "cf_access_aud", "") or "").strip()
119         if not team or not aud:
120             missing = " + ".join(n for n, v in (("cf_team_domain", team), ("cf_access_aud",
aud)) if not v)
121             out.append(StartupCheck(
122                 LEVEL_ERROR, "EXPOSED_EDGE_AUTH_INCOMPLETE",
123                 f"security.edge_auth_enabled is ON but {missing} is not configured ? the Cf-
Access JWT "
124                 "cannot be verified (every request would fail). Set security.cf_team_domain
"
125                 "(e.g. https://<team>.cloudflareaccess.com) and security.cf_access_aud (the
CF Access "
126                 "application AUD tag).",
127             ))
128
129         # (3) The auth extra (PyJWT[crypto]) is lazy-imported by
auth.verify_cf_access_token. If it's
130         # missing, every request 500s at verify time ? a boot WARN beats a per-request 500.
Warning, not
131         # fatal: the auth module's own ImportError is the final authority.
132         if not _auth_extra_available():
133             out.append(StartupCheck(
134                 LEVEL_WARN, "EXPOSED_AUTH_EXTRA_MISSING",
135                 "security.edge_auth_enabled is ON but PyJWT[crypto] is not importable ? Cf-
Access "
136                 "verification will fail at request time. Install the 'auth' extra (pip
install -e .[auth]).",
137             ))
138
139         return out
140
141

```

252. user (2026-06-21T15:31:00.943Z)

```
34:from __future__ import annotations
36:import logging
37:import os
38:from dataclasses import dataclass
39:from typing import Any, Mapping, Optional
41:LEVEL_ERROR = "error"
42:LEVEL_WARN = "warn"
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

253. user (2026-06-21T15:31:12.903Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\config\startup.py has been updated successfully. (file state is current in your context ? no need to Read it back)

254. assistant (2026-06-21T15:31:21.690Z)

Now refine the IPv6 dead entries (Starlette can't match them) and add a test for the new startup check. Let me find the existing exposure-check tests.

255. user (2026-06-21T15:31:24.421Z)

```
tests/test_bind_and_host.py
tests/test_edge_auth.py
tests/test_redact.py
tests/test_startup_checks.py
tests/__pycache__/test_bind_and_host.cpython-313-pytest-9.0.3.pyc
tests/__pycache__/test_edge_auth.cpython-313-pytest-9.0.3.pyc
tests/__pycache__/test_redact.cpython-313-pytest-9.0.3.pyc
tests/__pycache__/test_startup_checks.cpython-313-pytest-9.0.3.pyc
---
55:def test_no_profile_is_strict_noop(backend, device, detector):
61:def test_enforce_no_profile_never_raises_and_returns_empty():
68:def test_truly_local_local_storage_gpu_present_is_clean():
73:def test_truly_local_cuda_strict_without_gpu_warns():
80:def test_truly_local_auto_without_gpu_warns_cpu_fallback():
86:def test_truly_local_unprobed_cuda_skips_gpu_checks():
92:def test_truly_local_cloud_storage_warns_unusual():
98:def test_truly_local_gdrive_storage_is_fine():
105:def test_cloud_serving_local_storage_is_fatal():
112:def test_cloud_serving_gcs_without_creds_warns():
119:def test_cloud_serving_gcs_with_creds_is_clean():
125:def test_cloud_serving_on_cloud_run_env_skips_creds_warn():
131:def test_cloud_serving_cpu_device_warns():
137:def test_cloud_serving_no_gpu_warns():
145:def test_cpu_mock_stub_is_clean():
150:def test_cpu_mock_real_detector_warns():
158:def test_enforce_raises_on_fatal_and_carries_failures(caplog):
167:def test_enforce_does_not_raise_on_warnings_only():
173:def test_enforce_raise_on_error_false_returns_errors_without_raising():
182:def test_worker_cmd_infer_aborts_rc2_on_incoherent_profile(tmp_path):
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

256. user (2026-06-21T15:31:32.308Z)

```
18:    run_startup_checks,
22:def _cfg(profile="", backend="local", device="cuda", detector_impl="auto"):
36:    sec = {"edge_auth_enabled": edge_auth, "cf_team_domain": team, "cf_access_aud": aud}
38:    sec["allowed_video_root"] = video_root
46:def _codes(checks):
58:    assert run_startup_checks(cfg, cuda_available=False, env={}) == []
70:    assert run_startup_checks(cfg, cuda_available=True, env={}) == []
75:    checks = run_startup_checks(cfg, cuda_available=False, env={})
```

```

76:     assert _codes(checks) == {"CUDA_NOT_VISIBLE"}
82:     checks = run_startup_checks(cfg, cuda_available=False, env={})
83:     assert _codes(checks) == {"CPU_FALLBACK"}
89:     assert run_startup_checks(cfg, cuda_available=None, env={}) == []
94:     checks = run_startup_checks(cfg, cuda_available=True, env={})
95:     assert _codes(checks) == {"STORAGE_UNUSUAL"} # exactly this ? no spurious extra warning
100:     assert run_startup_checks(cfg, cuda_available=True, env={}) == []
107:     checks = run_startup_checks(cfg, cuda_available=True,
env={"GOOGLE_APPLICATION_CREDENTIALS": "/k"})
114:     checks = run_startup_checks(cfg, cuda_available=True, env={})
115:     assert _codes(checks) == {"CREDS_UNCONFIRMED"}
121:     assert run_startup_checks(cfg, cuda_available=True,
128:     assert run_startup_checks(cfg, cuda_available=True, env={"K_SERVICE": "rally"}) == []
133:     checks = run_startup_checks(cfg, cuda_available=True, env={"K_SERVICE": "rally"})
134:     assert _codes(checks) == {"CLOUD_CPU_DEVICE"}
139:     checks = run_startup_checks(cfg, cuda_available=False, env={"K_SERVICE": "rally"})
140:     assert _codes(checks) == {"CLOUD_CUDA_NOT_VISIBLE"}
147:     assert run_startup_checks(cfg, cuda_available=False, env={}) == []
152:     checks = run_startup_checks(cfg, cuda_available=False, env={})
153:     assert _codes(checks) == {"MOCK_NONSTUB_DETECTOR"}
198: def test_worker_default_off_skips_cuda_probe(monkeypatch):
204:     monkeypatch.setattr(wmain, "probe_cuda_available", lambda: (calls.append(1), False)[1])
205:     assert wmain._run_startup_checks(wmain.load_config()) is True # default config,
profile=""
209: def test_worker_active_profile_probes_cuda(monkeypatch, tmp_path):
214:     monkeypatch.setattr(wmain, "probe_cuda_available", lambda: (calls.append(1), True)[1])
217:     wmain._run_startup_checks(wmain.load_config(str(cfg_file)))
246:     """edge_auth_enabled=False (today's default) ? ZERO exposure checks, even with
include_exposure
247:     and an otherwise-unsafe config (no allowed_video_root). The strict default-OFF
guarantee."""
249:     assert run_startup_checks(cfg, include_exposure=True, env={}) == []
252: def test_exposure_not_run_without_flag(monkeypatch):
255:     monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
257:     assert run_startup_checks(cfg, env={}) == [] # include_exposure defaults False
260: def test_exposed_without_video_root_is_fatal(monkeypatch):
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

257. assistant (2026-06-21T15:31:42.405Z)

`run\_startup\_checks` takes an `env` param threaded to checks ? but my new check reads `os.environ` directly, which is inconsistent and hard to test. Let me see how `env` is threaded so I align with the existing pattern.

258. user (2026-06-21T15:31:43.369Z)

```

146             "defaults to loopback-only and will 400 forwarded requests. Set
RALLY_ALLOWED_HOSTS to your "
147             "public host(s) (or '*' when the edge gate validates the host).",
148         ))
149
150     return out
151
152
153     def run_startup_checks(
154         config: Any,
155         *,
156         cuda_available: Optional[bool] = None,
157         env: Optional[Mapping[str, str]] = None,
158         include_exposure: bool = False,
159     ) -> list[StartupCheck]:
160         """Return the startup findings for ``config``.
161
162         ``cuda_available`` is the caller's best-effort GPU probe (``None`` = not probed ?
GPU checks

```

```

163         are skipped; the API server passes ``None``, the worker passes
``presets.probe_cuda_available()``).
164         ``include_exposure`` adds the exposure-safety posture (``_exposure_checks``) ? the
HTTP **server**
165         passes ``True``; the worker leaves it ``False`` (it is never "exposed"). Both
dimensions are
166         DEFAULT-OFF: no profile ? no profile checks; ``edge_auth_enabled=False`` ? no
exposure checks.
167         Pure: no IO (beyond the ``find_spec`` auth probe), no torch import ? the caller owns
the GPU probe.
168         ""
169         e = os.environ if env is None else env
170         checks: list[StartupCheck] = []
171
172         # Exposure-safety posture ? INDEPENDENT of deployment.profile (a tunnel-to-a-box
exposure often
173         # runs with no profile set), gated on edge_auth_enabled (default-OFF). Server-only.
174         if include_exposure:
175             checks.extend(_exposure_checks(config))
176
177         profile = getattr(getattr(config, "deployment", None), "profile", "") or ""
178         if not profile:
179             return checks # no profile ? only the exposure posture (itself off unless edge
auth is on).
180
181         storage_backend = getattr(getattr(config, "storage", None), "backend", "local")
182         indexing = getattr(config, "indexing", None)
183         detector_impl = getattr(indexing, "detector_impl", "auto")
184         device = getattr(getattr(indexing, "wasb", None), "device", "cuda")
185
186         if profile == "truly-local":
187             # truly-local reads bytes from the local FS or a *mounted* Drive ? a cloud
backend here is
188             # an inconsistent (but not fatal) choice worth surfacing.
189             if storage_backend not in ("local", "gdrive"):
190                 checks.append(StartupCheck(
191                     LEVEL_WARN, "STORAGE_UNUSUAL",
192                     f"profile 'truly-local' usually uses local/gdrive storage, but
storage_backend="
193                     f"'{storage_backend}'. Reads will go to a remote backend.",
194                     ))
195             # GPU heads-up (warning only ? the detector healthcheck is the authority; the
probe is
196             # best-effort and may be False even when a subprocess detector env has CUDA).
197             if cuda_available is False:
198                 if device == "cuda":
199                     checks.append(StartupCheck(
200                         LEVEL_WARN, "CUDA_NOT_VISIBLE",
201                         "device='cuda' but no CUDA GPU is visible to this process. If the
box truly "
202                         "has no GPU the detector healthcheck will hard-fail ? set
indexing.wasb.device="
203                         "'auto' for the CPU fallback (M4).",
204                         ))
205                 elif device == "auto":

```

259. assistant (2026-06-21T15:31:59.697Z)

`env` is resolved in `run\_startup\_checks` but `\_exposure\_checks` only gets `config`. I'll thread
`env` into it so my check reads the passed env (testable + consistent), not `os.environ`
directly.

260. user (2026-06-21T15:32:04.948Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\config\startup.py has been updated

successfully. (file state is current in your context ? no need to Read it back)

261. user (2026-06-21T15:32:09.450Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\config\startup.py has been updated successfully. (file state is current in your context ? no need to Read it back)

262. user (2026-06-21T15:32:13.980Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\config\startup.py has been updated successfully. (file state is current in your context ? no need to Read it back)

263. assistant (2026-06-21T15:32:20.696Z)

Now refine the IPv6 dead entries in the host allowlist (Starlette's parser can't match `::1`).

264. user (2026-06-21T15:32:24.787Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

265. assistant (2026-06-21T15:32:32.320Z)

Now add a test for the new startup check. Let me read the exposure-test helper + an existing exposure test to mirror.

266. user (2026-06-21T15:32:33.344Z)

```
22     def _cfg(profile="", backend="local", device="cuda", detector_impl="auto"):
23         """Build an AppConfig with arbitrary profile/storage/device combos ? including ones
a
24         preset would never produce (e.g. cloud-serving + local storage), to exercise the
coherence
25         checks. model_validate fills every other field with its default."""
26         return AppConfig.model_validate({
27             "deployment": {"profile": profile},
28             "storage": {"backend": backend},
29             "indexing": {"detector_impl": detector_impl, "wasb": {"device": device}},
30         })
31
32
33     def _exposed_cfg(*, edge_auth=True, video_root="/jail",
team="https://acme.cloudflareaccess.com",
34             aud="abc123", profile="", backend="local"):
35         """An AppConfig with the security/exposure fields set, for the exposure-posture
checks."""
36         sec = {"edge_auth_enabled": edge_auth, "cf_team_domain": team, "cf_access_aud": aud}
37         if video_root is not None:
38             sec["allowed_video_root"] = video_root
39         return AppConfig.model_validate({
40             "deployment": {"profile": profile},
41             "storage": {"backend": backend},
42             "security": sec,
43         })
44
45
46     def _codes(checks):
47         return {c.code for c in checks}
48
49
50     # ----- default-OFF
51
52     @pytest.mark.parametrize("backend,device,detector", [
53         ("local", "cuda", "auto"), ("gcs", "cpu", "stub"), ("s3", "auto", "native"),
```

```

54     ])
55     def test_no_profile_is_strict_noop(backend, device, detector):
56         """profile='' ? ZERO checks regardless of any other field (the pre-M5 behaviour)."""
57         cfg = _cfg(profile="", backend=backend, device=device, detector_impl=detector)
58         assert run_startup_checks(cfg, cuda_available=False, env={}) == []
59
60
61     def test_enforce_no_profile_never_raises_and_returns_empty():
62         cfg = _cfg(profile="", backend="local")
63         assert enforce_startup_checks(cfg, cuda_available=False, env={}) == []
64
65
66     # ----- truly-local
67
68     def test_truly_local_local_storage_gpu_present_is_clean():
69         cfg = _cfg(profile="truly-local", backend="local", device="cuda")
70         assert run_startup_checks(cfg, cuda_available=True, env={}) == []
71
72
73     def test_truly_local_cuda_strict_without_gpu_warns():

```

267. assistant (2026-06-21T15:32:41.339Z)

My new check fires whenever edge-auth is on without `RALLY\_ALLOWED\_HOSTS`, which will affect existing exposure tests that assert exact code sets. Let me read them all.

268. user (2026-06-21T15:32:41.982Z)

```

242
243     # ----- exposure-safety posture (M9 hardening,
server-only)
244
245     def test_exposure_off_by_default_is_noop():
246         """edge_auth_enabled=False (today's default) ? ZERO exposure checks, even with
include_exposure
247         and an otherwise-unsafe config (no allowed_video_root). The strict default-OFF
guarantee."""
248         cfg = _exposed_cfg(edge_auth=False, video_root=None, team="", aud="")
249         assert run_startup_checks(cfg, include_exposure=True, env={}) == []
250
251
252     def test_exposure_not_run_without_flag(monkeypatch):
253         """The worker path (include_exposure=False, the default) NEVER runs exposure checks
? even with
254         edge auth ON and an unsafe config. Only the HTTP server opts in."""
255         monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
256         cfg = _exposed_cfg(edge_auth=True, video_root=None, team="", aud="")
257         assert run_startup_checks(cfg, env={}) == [] # include_exposure defaults False
258
259
260     def test_exposed_without_video_root_is_fatal(monkeypatch):
261         monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
262         cfg = _exposed_cfg(video_root=None) # team+aud set, jail missing
263         checks = run_startup_checks(cfg, include_exposure=True, env={})
264         assert _codes(checks) == {"EXPOSED_NO_VIDEO_ROOT"}
265         assert checks[0].level == LEVEL_ERROR
266
267
268     @pytest.mark.parametrize("video_root", [None, "", " "])
269     def test_exposed_blank_video_root_is_fatal(monkeypatch, video_root):
270         """None, empty, and whitespace-only roots are all treated as 'unset' ? fatal."""
271         monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
272         cfg = _exposed_cfg(video_root=video_root)
273         assert "EXPOSED_NO_VIDEO_ROOT" in _codes(run_startup_checks(cfg,
include_exposure=True, env={}))

```

```

274
275
276 @pytest.mark.parametrize("team,aud,expect_missing", [
277     ("", "abc", "cf_team_domain"),
278     ("https://t.cloudflareaccess.com", "", "cf_access_aud"),
279     ("", "", "cf_team_domain + cf_access_aud"),
280 ])
281 def test_exposed_incomplete_cf_config_is_fatal(monkeypatch, team, aud, expect_missing):
282     monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
283     cfg = _exposed_cfg(team=team, aud=aud) # jail set, CF config partial
284     checks = run_startup_checks(cfg, include_exposure=True, env={})
285     assert _codes(checks) == {"EXPOSED_EDGE_AUTH_INCOMPLETE"}
286     assert checks[0].level == LEVEL_ERROR
287     assert expect_missing in checks[0].message # names exactly what's missing
288
289
290 def test_exposed_missing_auth_extra_warns(monkeypatch):
291     """PyJWT not importable ? a WARN (never fatal ? the auth module's ImportError is the
292     authority)."""
293     monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: False)
294     cfg = _exposed_cfg() # fully configured otherwise
295     checks = run_startup_checks(cfg, include_exposure=True, env={})
296     assert _codes(checks) == {"EXPOSED_AUTH_EXTRA_MISSING"}
297     assert checks[0].level == LEVEL_WARN
298
299 def test_exposed_fully_configured_is_clean(monkeypatch):
300     monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
301     cfg = _exposed_cfg() # edge auth on, jail set, team+aud set, auth extra present

```

269. assistant (2026-06-21T15:33:06.487Z)

My new check fires under edge-auth without `RALLY\_ALLOWED\_HOSTS`, which collides with 4 existing exact-match exposure tests. That's expected ? I'll set `RALLY\_ALLOWED\_HOSTS` in those (orthogonal to what they test) and add a dedicated test for the new check. Let me see the "fully configured clean" test body.

270. user (2026-06-21T15:33:07.277Z)

```

299 def test_exposed_fully_configured_is_clean(monkeypatch):
300     monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
301     cfg = _exposed_cfg() # edge auth on, jail set, team+aud set, auth extra present
302     assert run_startup_checks(cfg, include_exposure=True, env={}) == []
303
304
305 def test_exposed_both_errors_together(monkeypatch):
306     """An unsafe exposure trips BOTH fatal checks at once (jail missing + CF config
307     missing)."""
308     monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
309     cfg = _exposed_cfg(video_root=None, team="", aud="")
310     codes = _codes(run_startup_checks(cfg, include_exposure=True, env={}))
311     assert codes == {"EXPOSED_NO_VIDEO_ROOT", "EXPOSED_EDGE_AUTH_INCOMPLETE"}

```

271. assistant (2026-06-21T15:33:30.997Z)

`test\_exposed\_both\_errors\_together` (line 310) is a 5th collision. I'll set `RALLY\_ALLOWED\_HOSTS` in the 5 exact-match exposure tests (orthogonal to their dimension) and add a dedicated test for the new check.

272. user (2026-06-21T15:33:33.854Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_startup_checks.py has been updated successfully. (file state is current in your context ? no need to Read it back)

273. user (2026-06-21T15:33:39.434Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_startup_checks.py has been updated successfully. (file state is current in your context ? no need to Read it back)

274. user (2026-06-21T15:33:45.314Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_startup_checks.py has been updated successfully. (file state is current in your context ? no need to Read it back)

275. user (2026-06-21T15:33:53.282Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_startup_checks.py has been updated successfully. (file state is current in your context ? no need to Read it back)

276. user (2026-06-21T15:34:02.313Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_startup_checks.py has been updated successfully. (file state is current in your context ? no need to Read it back)

277. assistant (2026-06-21T15:34:06.826Z)

Now run the affected tests, then the full suite.

278. user (2026-06-21T15:34:17.143Z)

```

                                env={"K_SERVICE": "x"}))
>      assert codes == {"EXPOSED_NO_VIDEO_ROOT", "STORAGE_INCOHERENT"}
E      AssertionError: assert {'EXPOSED_HOS...E_INCOHERENT'} == {'EXPOSED_NO...E_INCOHERENT'}
E
E      Extra items in the left set:
E      'EXPOSED_HOST_ALLOWLIST_LOOPBACK'
E      Use -v to get more diff

```

```

tests\test_startup_checks.py:331: AssertionError
===== short test summary info =====
FAILED tests/test_startup_checks.py::test_exposure_composes_with_profile_checks
1 failed, 77 passed in 6.62s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

279. assistant (2026-06-21T15:34:21.803Z)

One more exact-match test (line 331). Let me read and fix it.

280. user (2026-06-21T15:34:22.647Z)

```

322
323
324     def test_exposure_composes_with_profile_checks(monkeypatch):
325         """Exposure posture is INDEPENDENT of deployment.profile ? an exposed cloud-serving
config with
326         local storage trips both the exposure jail error AND the profile STORAGE_INCOHERENT
error."""
327         monkeypatch.setattr(startup_mod, "_auth_extra_available", lambda: True)
328         cfg = _exposed_cfg(video_root=None, profile="cloud-serving", backend="local")
329         codes = _codes(run_startup_checks(cfg, cuda_available=True, include_exposure=True,
330                                env={"K_SERVICE": "x"}))
331         assert codes == {"EXPOSED_NO_VIDEO_ROOT", "STORAGE_INCOHERENT"}
332
333

```

281. user (2026-06-21T15:34:28.128Z)

The file C:\Users\avidu\Projects\bhi-wt-p0c\tests\test_startup_checks.py has been updated successfully. (file state is current in your context ? no need to Read it back)

282. user (2026-06-21T15:34:45.251Z)

```
..... [100%]
55 passed in 5.71s
=== ruff+mypy ===
All checks passed!
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

283. assistant (2026-06-21T15:34:48.142Z)

Now the full suite + coverage gate.

284. user (2026-06-21T15:36:46.756Z)

```
backend\config\startup.py      86      3    97%
backend\main.py                445     130   71%
TOTAL                          9695   1442   85%
Required test coverage of 70% reached. Total coverage: 85.13%
1271 passed, 7 skipped, 7 warnings in 113.83s (0:01:53)
RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

285. assistant (2026-06-21T15:36:53.714Z)

Full suite green (1271 passed), coverage 85.13%, startup.py 97%. All review findings folded.
Committing P0c + opening the PR.

286. user (2026-06-21T15:37:11.283Z)

```
=== status (only P0c files) ===
 M backend/config/startup.py
 M backend/main.py
 M tests/test_startup_checks.py
?? tests/test_bind_and_host.py
warning: in the working copy of 'tests/test_bind_and_host.py', LF will be replaced by CRLF the
next time Git touches it
=== HEAD vs origin/master ===
6b78c0a feat(pkg): P0c ? config/profile-driven bind address + DNS-rebinding Host-header guard
(default-OFF)
branch 'feat/packaging-p0c-bind-host' set up to track 'origin/feat/packaging-p0c-bind-host'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
 * [new branch]      feat/packaging-p0c-bind-host -> feat/packaging-p0c-bind-host
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

287. user (2026-06-21T15:37:28.723Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/296>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

288. assistant (2026-06-21T15:37:31.855Z)

PR #296 opened: [feat(pkg): P0c ? bind + Host-header guard](<https://github.com/Khelsutra/badminton-highlight-indexer/pull/296>). Watching CI.

289. user (2026-06-21T15:39:57.998Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    23s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187758
```

```
Full test suite (offline, mocked)      pass    2m22s  https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187744
Golden fixtures cross-OS (macos-latest) pass    16s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187765
Golden fixtures cross-OS (ubuntu-latest) pass    21s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187786
Golden fixtures cross-OS (windows-latest) pass    1m1s  https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187766
Mypy (lenient baseline) pass    30s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187784
Ruff lint      pass    11s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187745
Syntax sweep + import smoke (no GPU stack) pass    23s  https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27909175964/job/82583187758
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

290. assistant (2026-06-21T15:40:01.326Z)

All 7 CI green. Merging P0c and cleaning up.

291. user (2026-06-21T15:40:10.056Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
fc78fce..3895963 master -> origin/master
=== master ===
3895963 feat(pkg): P0c ? config/profile-driven bind address + DNS-rebinding Host-header guard
(default-OFF) (#296)
=== PR #296 ===
MERGED
```

292. assistant (2026-06-21T15:40:22.446Z)

P0c ****merged**** (master `3895963`). 3 of 5 P0 PRs done. Updating tracker + memory, then P0d.

293. user (2026-06-21T15:40:30.254Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

294. user (2026-06-21T15:40:43.161Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

295. assistant (2026-06-21T15:41:10.202Z)

Now ****P0d ? declare `scipy` + license-scan & built-artifact CI lanes**** (independent of `main.py`). Let me set up the worktree and study the CI structure so the new lanes are deterministic.

296. user (2026-06-21T15:41:15.840Z)

```
HEAD is now at 3895963 feat(pkg): P0c ? config/profile-driven bind address + DNS-rebinding Host-header guard (default-OFF) (#296)
=== scipy declared anywhere? ===
(scipy NOT declared)
=== installed scipy version ===
scipy 1.17.1
```

```

=== ci.yml job names + structure ===
3:on:
4:  push:
6:  pull_request:
8:permissions:
11:concurrency:
15:jobs:
20:  lint:
23:    steps:
26:      with:
29:      run: python -m pip install ruff
37:  typecheck:
40:    steps:
43:      with:
46:      run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-genai
numpy
54:  import-smoke:
55:    name: Syntax sweep + import smoke (no GPU stack)
57:    steps:
60:      with:
63:      run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-genai
64:    - name: Run import smoke tests
65:      run: python -m pytest tests/test_import_smoke.py -v
67:  full-suite:
70:    steps:
73:      with:
84:      run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
88:      run: python -m pip install -e .[dev]
97:      run: python run_all_tests.py --cov=backend --cov=training --cov-report=term --cov-
fail-under=70
105:  golden-crossos:
106:    name: Golden fixtures cross-OS (${{ matrix.os }})
107:    strategy:
109:      matrix:
110:        os: [ubuntu-latest, windows-latest, macos-latest]
111:      runs-on: ${{ matrix.os }}
112:    steps:
115:      with:
118:      run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
124:      python -m backend.eval.golden_fixtures --check
125:      python -m pytest tests/test_golden_fixtures.py tests/test_golden_real_fixtures.py
-q

```

297. assistant (2026-06-21T15:41:19.671Z)

Let me read the full ci.yml to mirror its exact format for the new lanes.

298. user (2026-06-21T15:41:20.444Z)

```

1      name: CI
2
3      on:
4        push:
5          branches: [master]
6        pull_request:
7
8      permissions:
9        contents: read
10
11     concurrency:
12       group: ci-${{ github.ref }}
13       cancel-in-progress: true
14

```

```

15 jobs:
16   # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17   # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
18   # tracked .py, including scripts/ and training/, which the backend-only
19   # py_compile smoke sweep does not reach.
20   lint:
21     name: Ruff lint
22     runs-on: ubuntu-latest
23     steps:
24       - uses: actions/checkout@v6
25       - uses: actions/setup-python@v6
26         with:
27           python-version: "3.13"
28       - name: Install ruff
29         run: python -m pip install ruff
30       - name: Ruff check
31         run: ruff check . --output-format github
32
33   # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
34   # per-module quarantine list ? ratchet by removing entries as modules clean
35   # up. Typed core deps are installed so mypy sees their real signatures
36   # (without them ignore_missing_imports would Any-stub half the checks).
37   typecheck:
38     name: Mypy (lenient baseline)
39     runs-on: ubuntu-latest
40     steps:
41       - uses: actions/checkout@v6
42       - uses: actions/setup-python@v6
43         with:
44           python-version: "3.13"
45       - name: Install mypy + typed core deps
46         run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
genai numpy
47       - name: Mypy
48         run: mypy backend training
49
50   # Fast guard against the 7fbb0 class of breakage: a SyntaxError or
51   # module-level import error shipping while the suite stays green.
52   # Deliberately does NOT install cv2/torch/numpy ? the smoke tests must hold
53   # on machines without the GPU stack (test_import_smoke stubs what's absent).
54   import-smoke:
55     name: Syntax sweep + import smoke (no GPU stack)
56     runs-on: ubuntu-latest
57     steps:
58       - uses: actions/checkout@v6
59       - uses: actions/setup-python@v6
60         with:
61           python-version: "3.13"
62       - name: Install minimal deps
63         run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai
64       - name: Run import smoke tests
65         run: python -m pytest tests/test_import_smoke.py -v
66
67   full-suite:
68     name: Full test suite (offline, mocked)
69     runs-on: ubuntu-latest
70     steps:
71       - uses: actions/checkout@v6
72       - uses: actions/setup-python@v6
73         with:
74           python-version: "3.13"
75       - cache: pip
76       - name: System libs for opencv-python
77         run: |

```

```

78         sudo apt-get update
79         sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
80     - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
81       # torch/torchvision are NOT in pyproject deps ? their CPU/CUDA wheels
82       # need an out-of-band index that PEP 621 can't express. Keep installing
83       # them separately here, before the editable install pulls the rest.
84       run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
85     - name: Install package (editable) + dev tooling
86       # Editable PEP 621 install: runtime deps + the `dev` group
87       # (pytest, pytest-cov, httpx for the fastapi TestClient, ruff, mypy).
88       run: python -m pip install -e .[dev]
89       # obsreport (private, soft dep) is absent here: reporting tests skip via
90       # pytest.importorskip ? run_all_tests.py prints which. The coverage floor
91       # is therefore calibrated against THIS lane's number (local runs with
92       # obsreport installed measure a little higher).
93       # Floor calibrated 2026-06-11 against this lane's measured 72% (local
94       # with obsreport: 73%). Ratchet it UP as coverage grows; never down
95       # without an owner decision.
96     - name: Run full suite (with coverage floor)
97       run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
98
99     # Cross-OS determinism guard for the committed golden regression fixtures
100    # (docs/GOLDEN_REGRESSION_FIXTURES.md). The CONTROL replay path is pure-stdlib
101    # (only numpy, no cv2/torch), so this is a FAST, focused job ? it does NOT run
102    # the full suite x3 OSes. It replays the Windows/Linux/macOS runner against the
103    # committed (parse-compared) snapshots, catching any cross-OS float drift beyond
104    # the 1e-6 tolerance before it can flake a contributor's PR on another machine.
105    golden-crossos:
106      name: Golden fixtures cross-OS (${{ matrix.os }})
107      strategy:
108        fail-fast: false
109        matrix:
110          os: [ubuntu-latest, windows-latest, macos-latest]
111      runs-on: ${{ matrix.os }}
112      steps:
113        - uses: actions/checkout@v6
114        - uses: actions/setup-python@v6
115          with:
116            python-version: "3.13"
117        - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
118          run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
119        - name: Golden-fixtures characterization (cross-OS determinism guard)
120          # Also runs the freeze-real / time-origin-reset / M2 refusal-gate tests
(synthetic /
121          # tmp_path only ? no GPU/network) so the reset metric-invariance (fix #2) and
the
122          # de-attribution paths get cross-OS coverage and surface any win/mac float
drift.
123          run: |
124            python -m backend.eval.golden_fixtures --check
125            python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
126

```

299. assistant (2026-06-21T15:42:08.250Z)

CI structure is clear. Before designing the license-scan lane (which must not flake), let me run `pip-licenses` over the dependency closure now to see what license strings actually appear ? verify-driven, so my denylist is accurate.

300. user (2026-06-21T15:42:43.358Z)

```
[notice] A new release of pip is available: 26.0.1 -> 26.1.2
[notice] To update, run: python.exe -m pip install --upgrade pip
=== licenses of the RUNTIME deps + key transitive ===
Name          Version      License
Jinja2         3.1.6        BSD License
anyio          4.13.0        MIT
click          8.4.0        BSD-3-Clause
fastapi        0.136.1       MIT
google-genai   2.4.0        Apache-2.0
h11            0.16.0        MIT License
numpy          2.4.6        BSD-3-Clause AND 0BSD AND MIT AND Zlib AND CC0-1.0
opencv-python  4.13.0.92     Apache Software License
pydantic       2.13.4        MIT
python-dotenv  1.2.2        BSD-3-Clause
scipy          1.17.1        BSD License
starlette      1.0.0        BSD-3-Clause
supervision    0.28.0        MIT
uvicorn        0.47.0        BSD-3-Clause
=== any copyleft (AGPL/SSPL/GPL/LGPL) ANYWHERE in the full env? ===
PyMuPDF              1.27.2.3      Dual Licensed - GNU AFFERO GPL 3.0 or Artifex
Commercial License
ultralitics          8.4.51        GNU Affero General Public License v3 or later (AGPLv3+)
ultralitics-thop     2.0.19        GNU Affero General Public License v3 or later (AGPLv3+)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

301. assistant (2026-06-21T15:43:57.771Z)

Verified ? **every runtime dep + transitive is permissive** (MIT/Apache/BSD/0BSD/Zlib/CC0); the only AGPL hits (`ultralitics`, `PyMuPDF`) are NOT in the `pip install .` closure (dev-env noise; ultralytics was deliberately removed). So a copyleft-denylist license lane will pass deterministically. Implementing P0d. First, declare scipy.

302. user (2026-06-21T15:43:58.045Z)

```
1      fastapi>=0.111.0
2      uvicorn>=0.30.1
3      opencv-python>=4.9.0.80
4      numpy>=1.26.4
5      pydantic>=2.7.1
6      jinja2>=3.1.4
7      python-dotenv>=1.0.0
8      google-genai>=2.4.0
9      torch>=2.12.0
10     torchvision>=0.27.0
11     supervision>=0.21.0,<0.30.0
12
13     # Per-video observability reports (SOFT dependency: the pipeline degrades to a
14     # no-op if absent). Shared, permissively-licensed report toolkit. Pin a tag once
15     # the private repo is published, e.g.:
16     #   obsreport @ git+ssh://git@github.com/avidullu/sports-obsreport.git@v0.1.3
17     # For local development install it editable from a sibling checkout:
18     #   pip install -e ../sports-obsreport --no-build-isolation
19
```

303. user (2026-06-21T15:44:04.064Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

304. user (2026-06-21T15:44:08.615Z)

```
26     # Runtime deps mirror requirements.txt, EXCEPT torch/torchvision (see header).
27     dependencies = [
28         "fastapi>=0.111.0",
29         "uvicorn>=0.30.1",
```

```

30         "opencv-python>=4.9.0.80",
31         "numpy>=1.26.4",
32         "pydantic>=2.7.1",
33         "jinja2>=3.1.4",
34         "python-dotenv>=1.0.0",
35         "google-genai>=2.4.0",
36         "supervision>=0.21.0,<0.30.0",
37     ]

```

305. user (2026-06-21T15:44:13.934Z)

The file C:\Users\avidu\Projects\bhi-wt-p0d\pyproject.toml has been updated successfully. (file state is current in your context ? no need to Read it back)

306. user (2026-06-21T15:44:18.604Z)

The file C:\Users\avidu\Projects\bhi-wt-p0d\requirements.txt has been updated successfully. (file state is current in your context ? no need to Read it back)

307. assistant (2026-06-21T15:44:37.246Z)

To validate before adding CI lanes (verify-driven), let me build the wheel and confirm it includes `backend`+`training`+data, and that `indexer --help` launches from a non-repo dir. I'll build only (no install into the shared env, which would clobber the editable install other worktrees use).

308. user (2026-06-21T15:44:47.975Z)

```

[notice] To update, run: python.exe -m pip install --upgrade pip
=== build wheel ===
* Installing packages in isolated environment:
  - hatchling
* Getting build dependencies for wheel...
* Building wheel...
Successfully built badminton_highlight_indexer-0.1.0-py3-none-any.whl
=== wheel contents (backend/training/font present?) ===
files: 146
has backend/main.py: True
has training: True
has app_home: True
font: ['backend/assets/fonts/RobotoSlab-Regular.ttf']
=== indexer --help from /tmp (entrypoint + namespace import) ===
/usr/bin/bash: line 7: indexer: command not found
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

309. assistant (2026-06-21T15:45:01.059Z)

Wheel builds cleanly and **includes `backend/main.py`, `training/`, `app\_home.py`, and the bundled font** ? the packaging surface is complete. (`indexer` "not found" is just the `.exe` shim not being on Git-Bash's PATH locally; on CI's ubuntu it's on PATH.) Let me confirm the module entrypoint launches from a non-repo dir.

310. user (2026-06-21T15:45:03.078Z)

```

Exit code 1
=== python -m backend.main --help from /tmp ===
C:\Users\avidu\AppData\Local\Programs\Python\Python313\python.exe: Error while finding module
specification for 'backend.main' (ModuleNotFoundError: No module named 'backend')
...
exit: 0
=== entrypoint importable ===
Traceback (most recent call last):
  File "<string>", line 1, in <module>
    from backend.main import main; print('main callable:', callable(main))

```



```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
ModuleNotFoundError: No module named 'backend'
```

311. assistant (2026-06-21T15:45:38.595Z)

`backend` isn't globally importable in this dev env (the install is cwd-dependent) ? which is exactly the packaging gap this whole effort addresses, and means I can't smoke the wheel from `/tmp` here. Let me prove the lane properly in a throwaway venv (no torch needed to validate install + namespace resolution).

312. user (2026-06-21T15:47:12.687Z)

```
=== install ONLY the wheel (no torch) into a clean venv ===
```

```
[notice] A new release of pip is available: 26.0.1 -> 26.1.2
[notice] To update, run: C:\Users\avidu\AppData\Local\Temp\p0d_venv\Scripts\python.exe -m pip
install --upgrade pip
=== from a NON-repo cwd: does the wheel expose backend+training importably? ===
namespace import OK: None
=== does importing backend.main need torch? (tells the lane whether to install cpu torch) ===
File "C:\Users\avidu\AppData\Local\Temp\p0d_venv\Lib\site-
packages\backend\pipeline\detectors\person_detector.py", line 29, in <module>
    import torch
ModuleNotFoundError: No module named 'torch'
=== is the indexer console script installed? ===
indexer.exe
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

313. assistant (2026-06-21T15:47:21.965Z)

Wheel installs into a clean venv and exposes `backend`/`training` from any cwd; `import backend.main` needs torch (`person\_detector.py:29`), so the lane must install CPU torch (like the full-suite lane). Let me finish the proof: install CPU torch + run `indexer --help`.

314. user (2026-06-21T15:49:16.350Z)

```
=== install CPU torch into the venv ===
[notice] A new release of pip is available: 26.0.1 -> 26.1.2
[notice] To update, run: C:\Users\avidu\AppData\Local\Temp\p0d_venv\Scripts\python.exe -m pip
install --upgrade pip
=== indexer --help from /tmp (the exact built-artifact smoke) ===
usage: indexer [-h] [--video-path VIDEO_PATH] [--output-dir OUTPUT_DIR]
              [--setup] [--debug-clip DEBUG_CLIP] [--server] [--port PORT]
              [--host HOST] [--max-frames MAX_FRAMES]
              [--segmenter
{motion,gemini,yolo_hybrid,tracknet_hybrid,wasb_hybrid,fusion_hybrid}]
              [--trim-start TRIM_START] [--trim-end TRIM_END]
              [--trim-output TRIM_OUTPUT]

[exit 0]
=== and python -m backend.main --help ===
usage: main.py [-h] [--video-path VIDEO_PATH] [--output-dir OUTPUT_DIR]
              [--setup] [--debug-clip DEBUG_CLIP] [--server] [--port PORT]

[exit 0]
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

315. assistant (2026-06-21T15:49:51.362Z)

****Definitive proof**:** `indexer --help` runs (exit 0) from `/tmp` in a clean venv (wheel + CPU torch) ? and the help even shows the new `--host` flag, confirming consistency with master. The built-artifact lane commands are validated end-to-end. Cleaning up the throwaway venv/dist, then adding both CI lanes.

316. user (2026-06-21T15:50:01.398Z)

```
cleaned venv + dist
M pyproject.toml
M requirements.txt
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

317. assistant (2026-06-21T15:50:04.278Z)

Now append the two CI lanes (validated commands).

318. user (2026-06-21T15:50:15.704Z)

The file C:\Users\avidu\Projects\bhi-wt-p0d\github\workflows\ci.yml has been updated successfully. (file state is current in your context ? no need to Read it back)

319. assistant (2026-06-21T15:50:23.067Z)

Let me validate the YAML and run the full suite (the gate, even though P0d adds no Python code).

320. user (2026-06-21T15:50:25.591Z)

```
=== YAML valid? ===
jobs: ['lint', 'typecheck', 'import-smoke', 'full-suite', 'golden-crossos', 'license-scan',
'build-artifact']
=== pyproject parses + scipy present ===
['scipy>=1.10']
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

321. user (2026-06-21T15:51:59.295Z)

```
1271 passed, 7 skipped in 86.71s (0:01:26)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

322. assistant (2026-06-21T15:52:07.825Z)

Full suite green (1271), YAML valid (7 jobs). P0d is config/CI-only and I've validated the wheel-smoke end-to-end locally + confirmed the license closure is clean, so the new lanes will be exercised for real on this PR's CI. Committing.

323. user (2026-06-21T15:52:21.113Z)

```
=== status ===
M .github/workflows/ci.yml
M pyproject.toml
M requirements.txt
=== HEAD ===
86d023a feat(pkg): P0d ? declare scipy (undeclared serving dep) + license-scan & built-wheel CI
lanes
branch 'feat/packaging-p0d-scipy-ci' set up to track 'origin/feat/packaging-p0d-scipy-ci'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch]      feat/packaging-p0d-scipy-ci -> feat/packaging-p0d-scipy-ci
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

324. user (2026-06-21T15:52:36.874Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/297>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

325. assistant (2026-06-21T15:52:40.920Z)

****PR #297 opened: [feat(pkg): P0d ? scipy + CI lanes](<https://github.com/Khelsutra/badminton-highlight-indexer/pull/297>)**

highlight-indexer/pull/297).** This one exercises the two new lanes for the first time ?
watching closely.

326. user (2026-06-21T15:55:07.227Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    20s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223126
Built wheel smoke (install + launch)      pass    1m8s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27909563644/job/82584223149
Dependency license scan (no copleft in runtime deps)  pass    40s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223133
Full test suite (offline, mocked)      pass    2m22s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27909563644/job/82584223180
Golden fixtures cross-OS (macos-latest) pass    17s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27909563644/job/82584223154
Golden fixtures cross-OS (ubuntu-latest)      pass    20s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223128
Golden fixtures cross-OS (windows-latest)      pass    1m9s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223147
Mypy (lenient baseline) pass    26s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223156
Ruff lint      pass    13s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223104
Syntax sweep + import smoke (no GPU stack)      pass    20s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27909563644/job/82584223126
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

327. assistant (2026-06-21T15:55:11.562Z)

All 9 lanes green ? including both new ones (**Built wheel smoke** ?, **Dependency license scan** ?). Merging P0d.

328. user (2026-06-21T15:55:19.879Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:\Users\avidu\Projects\badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
3895963..e8bad40 master -> origin/master
=== master ===
e8bad40 feat(pkg): P0d ? declare scipy (undeclared serving dep) + license-scan & built-wheel CI
lanes (#297)
=== PR #297 ===
MERGED
```

329. assistant (2026-06-21T15:56:55.148Z)

P0d **merged** (master `e8bad40`). **4 of 5 P0 PRs done.** Final one: **P0e ? single-instance guard**. Setting up the worktree and reading the current server-init block (shifted by P0c).

330. user (2026-06-21T15:56:58.971Z)

```
HEAD is now at e8bad40 feat(pkg): P0d ? declare scipy (undeclared serving dep) + license-scan &
built-wheel CI lanes (#297)
=== head ===
e8bad40 feat(pkg): P0d ? declare scipy (undeclared serving dep) + license-scan & built-wheel CI
lanes (#297)
=== server-init: db_path + Database + fail_stale ===
720:    global global_config, db_instance
721:    global_config = load_app_config() # returns typed AppConfig
```

```

729: global_config.output.output_dir = resolve_output_dir(args.output_dir)
732: db_path = os.path.join(global_config.output.output_dir, global_config.output.db_name)
733: db_instance = Database(db_path)
737: swept = db_instance.fail_stale_jobs()
805: print(f"[SUCCESS] Process complete. Indexed {len(segments)} play segments
inside database: {db_path}")

```

331. user (2026-06-21T15:57:02.501Z)

```

718
719     # Initialize Global Config and DB
720     global global_config, db_instance
721     global_config = load_app_config() # returns typed AppConfig
722     _enforce_startup_or_exit(global_config) # M5: fail fast on an incoherent ACTIVE
profile
723
724     # The CLI --output-dir overrides the configured output directory. It now lives
725     # on the typed model (OutputConfig.output_dir), so every consumer ? including
726     # /api/compile ? reads the same value instead of a write-only runtime hack.
727     # P0a: a RELATIVE output dir roots under the app-home (RALLY_APP_HOME); with the env
728     # unset, resolve_output_dir("output") == "output" (CWD-relative) ? byte-identical to
before.
729     global_config.output.output_dir = resolve_output_dir(args.output_dir)
730     os.makedirs(global_config.output.output_dir, exist_ok=True)
731
732     db_path = os.path.join(global_config.output.output_dir,
global_config.output.db_name)
733     db_instance = Database(db_path)
734
735     # Crash recovery (#16): jobs left queued/running by a previous process are dead ?
736     # the executor is in-process. Mark them failed so pollers see a terminal state.
737     swept = db_instance.fail_stale_jobs()
738     if swept:
739         logger.warning(f"[JOBS] Swept {swept} stale job(s) from a previous run ?
failed.")
740
741     # CLI processing mode (no web UI needed)

```

332. assistant (2026-06-21T15:57:07.151Z)

Now write the cross-platform single-instance lock module (OS-managed lock, auto-released on crash ? no stale-PID problem).

333. assistant (2026-06-21T15:57:19.710Z)

""Single-instance file lock (PACKAGING_PLAN.md P0e).

The serving model runs an IN-PROCESS job worker, and ``fail\_stale\_jobs()`` sweeps any job left ``queued``/``running`` by a "previous" process to ``failed`` on startup. If a SECOND process opens the same SQLite DB (a second ``indexer --server``, or a CLI run while a server is up), its sweep would mark the FIRST process's in-flight jobs failed ? corrupting live state. This module guards the DB with an exclusive OS lock so the second instance refuses to start instead.

OS-managed advisory lock (``fcntl.flock`` on POSIX, ``msvcrt.locking`` on Windows): held for the process lifetime via an open file handle and **released automatically by the OS on exit OR crash**, so there is no stale-lock problem (unlike a bare PID file). The lock is keyed to the lock-file path, which callers co-locate with the DB ? so two instances on DIFFERENT ``--output-dir`` (different DBs) coexist fine.

```
"""
```

```
from __future__ import annotations
```

```
import os
```

```
from typing import IO, Optional
```

```
_IS_WINDOWS = os.name == "nt"
```

```
def acquire_lock(lock_path: str) -> Optional[IO]:
```

```
    """Try to take an exclusive, NON-blocking lock on ``lock_path``.
```

```
    Returns the open file handle on success (KEEP IT ALIVE for the process lifetime ? letting it
be garbage-collected/closed releases the lock), or ``None`` if another live process holds it.
    """
```

```
    parent = os.path.dirname(os.path.abspath(lock_path))
```

```
    if parent:
```

```
        os.makedirs(parent, exist_ok=True)
```

```
    fh = open(lock_path, "a+")
```

```
    try:
```

```
        if _IS_WINDOWS:
```

```
            import msvcrt
```

```
            fh.seek(0)
```

```
            msvcrt.locking(fh.fileno(), msvcrt.LK_NBLCK, 1) # lock 1 byte at offset 0, non-
```

```
blocking
```

```
        else:
```

```
            import fcntl
```

```
            fcntl.flock(fh.fileno(), fcntl.LOCK_EX | fcntl.LOCK_NB)
```

```
    except OSError:
```

```
        fh.close()
```

```
        return None
```

```
    # Record our PID for human diagnostics (best-effort; the lock ? not this content ? is
authoritative).
```

```
    try:
```

```
        fh.seek(0)
```

```
        fh.truncate()
```

```
        fh.write(str(os.getpid()))
```

```
        fh.flush()
```

```
    except OSError:
```

```
        pass
```

```
    return fh
```

```
def release_lock(fh: Optional[IO]) -> None:
```

```
    """Release a lock taken by :func:`acquire_lock` (the OS also releases it on process
exit)."""
```

```
    if fh is None:
```

```
        return
```

```
    try:
```

```
        if _IS_WINDOWS:
```

```
            import msvcrt
```

```
            fh.seek(0)
```

```
            msvcrt.locking(fh.fileno(), msvcrt.LK_UNLCK, 1)
```

```
        else:
```

```
            import fcntl
```

```
            fcntl.flock(fh.fileno(), fcntl.LOCK_UN)
```

```
    except OSError:
```

```
        pass
```

```
    finally:
```

```
        try:
```

```
            fh.close()
```

```
        except OSError:
```

```
            pass
```

334. user (2026-06-21T15:57:19.734Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0e\backend\utils\single_instance.py (file state is current in your context ? no need to Read it back)

335. assistant (2026-06-21T15:57:23.492Z)

Now wire it into `main.py` (lock the DB before the destructive `fail\_stale\_jobs` sweep) and add the import + module global.

336. user (2026-06-21T15:57:30.482Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

337. assistant (2026-06-21T15:57:35.446Z)

Add the import and the module-global for the lock handle.

338. user (2026-06-21T15:57:37.498Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

339. user (2026-06-21T15:57:41.230Z)

```
94         return hosts or ["localhost", "127.0.0.1", "testserver"]
95
96
97     # DNS-rebinding guard. Bound to localhost by default, so only loopback Host headers are
honoured
98     # (a remote page can't read responses cross-origin AND its Host header is rejected
here). www_redirect
99     # off so an unexpected ``www.`` host is rejected rather than 307-redirceted.
```

340. user (2026-06-21T15:57:46.141Z)

```
144:db_instance: Optional[Database] = None
145:global_config: Optional[AppConfig] = None # now typed (Pydantic model)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

341. user (2026-06-21T15:57:50.285Z)

```
143     # Global variables initialized at startup
144     db_instance: Optional[Database] = None
145     global_config: Optional[AppConfig] = None # now typed (Pydantic model)
146
```

342. user (2026-06-21T15:57:57.391Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

343. assistant (2026-06-21T15:58:04.769Z)

Now the unit tests for the lock mechanism.

344. assistant (2026-06-21T15:58:12.384Z)

```
"""Single-instance OS lock (PACKAGING_PLAN.md P0e)."""
```

```
import os
```

```

from backend.utils.single_instance import acquire_lock, release_lock

def test_acquire_then_second_acquire_is_blocked(tmp_path):
    lock = str(tmp_path / "db.sqlite.lock")
    fh = acquire_lock(lock)
    assert fh is not None, "first acquire should succeed"
    try:
        # A second acquire on the SAME lock path (even from this process, via a new handle) must
fail
        # while the first is held ? that's the second-instance-refused guarantee.
        assert acquire_lock(lock) is None
    finally:
        release_lock(fh)

def test_release_allows_reacquire(tmp_path):
    lock = str(tmp_path / "db.sqlite.lock")
    fh = acquire_lock(lock)
    assert fh is not None
    release_lock(fh)
    fh2 = acquire_lock(lock) # released ? free again
    assert fh2 is not None
    release_lock(fh2)

def test_different_paths_dont_conflict(tmp_path):
    a = acquire_lock(str(tmp_path / "a.lock"))
    b = acquire_lock(str(tmp_path / "b.lock")) # different DB/output-dir ? independent
    assert a is not None and b is not None
    release_lock(a)
    release_lock(b)

def test_acquire_creates_parent_dir(tmp_path):
    nested = tmp_path / "deep" / "nested" / "db.lock"
    fh = acquire_lock(str(nested))
    try:
        assert fh is not None
        assert os.path.exists(str(nested))
    finally:
        release_lock(fh)

def test_lockfile_records_pid(tmp_path):
    lock = str(tmp_path / "db.lock")
    fh = acquire_lock(lock)
    try:
        with open(lock, "r", encoding="utf-8") as f:
            assert f.read().strip() == str(os.getpid())
    finally:
        release_lock(fh)

```

345. user (2026-06-21T15:58:12.398Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p0e\tests\test_single_instance.py
(file state is current in your context ? no need to Read it back)

346. assistant (2026-06-21T15:58:26.446Z)

Now the subprocess integration test ? two servers on the same DB, the second must refuse.

347. user (2026-06-21T15:58:38.473Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

348. user (2026-06-21T15:58:44.069Z)

```
130         "DB was NOT created under RALLY_APP_HOME ? resolve_output_dir is not wired into
main() init.\n"
131         f"stdout:\n{proc.stdout[-1500:]} \nstderr:\n{proc.stderr[-1500:]}")
132     )
133     assert not (launch / "output").exists(), "output/ leaked into the launch CWD (should
be under the app home)"
134
```

349. user (2026-06-21T15:58:56.736Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\tests\test_server_entrypoint.py has been updated successfully. (file state is current in your context ? no need to Read it back)

350. assistant (2026-06-21T15:59:00.290Z)

Now run the P0e tests (unit + the two-server integration), then lint/type.

351. user (2026-06-21T15:59:31.349Z)

```
tmp_path = WindowsPath('C:/Users/avidu/AppData/Local/Temp/pytest-of-
avidu/pytest-786/test_lockfile_records_pid0')
```

```
def test_lockfile_records_pid(tmp_path):
    lock = str(tmp_path / "db.lock")
    fh = acquire_lock(lock)
    try:
        with open(lock, "r", encoding="utf-8") as f:
            assert f.read().strip() == str(os.getpid())
            ^^^^^^^^
E           PermissionError: [Errno 13] Permission denied
```

```
tests\test_single_instance.py:53: PermissionError
===== short test summary info =====
FAILED tests/test_single_instance.py::test_lockfile_records_pid - PermissionE...
1 failed, 7 passed in 26.36s
=== ruff+mypy ===
All checks passed!
backend\utils\single_instance.py:67: error: Module has no attribute "flock" [attr-defined]
backend\utils\single_instance.py:67: error: Module has no attribute "LOCK_UN" [attr-defined]
Found 5 errors in 1 file (checked 1 source file)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

352. assistant (2026-06-21T16:00:04.829Z)

Two cross-platform fixes: (1) use `sys.platform == "win32"` inline (mypy does platform-narrowing on `sys.platform`, not `os.name`, so the POSIX/Windows branches stop cross-erroring), and (2) the PID-readback test must read *after* release (the locked byte blocks reads on Windows). Rewriting the module.

353. user (2026-06-21T16:00:13.396Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\utils\single_instance.py has been updated successfully. (file state is current in your context ? no need to Read it back)

354. user (2026-06-21T16:00:19.258Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\utils\single_instance.py has been updated successfully. (file state is current in your context ? no need to Read it back)

355. user (2026-06-21T16:00:25.631Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\tests\test_single_instance.py has been updated successfully. (file state is current in your context ? no need to Read it back)

356. user (2026-06-21T16:00:32.507Z)

```
..... [100%]
5 passed in 0.12s
=== mypy ===
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

357. assistant (2026-06-21T16:00:39.085Z)

Unit tests green, mypy clean. Now the integration test (two real server boots ? the riskiest), then the full suite.

358. user (2026-06-21T16:01:07.307Z)

```
... [100%]
3 passed in 25.51s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

359. assistant (2026-06-21T16:01:11.626Z)

The two-server lock test passes. Full suite + coverage + final lint.

360. user (2026-06-21T16:03:22.177Z)

```
=== RUFF ===
All checks passed!
=== FULL SUITE + COV ===
backend\main.py 451 132 71%
backend\utils\single_instance.py 43 11 74%
TOTAL 9744 1455 85%
Required test coverage of 70% reached. Total coverage: 85.07%
1277 passed, 7 skipped, 7 warnings in 125.07s (0:02:05)
RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

361. assistant (2026-06-21T16:03:43.247Z)

Full suite green (1277 passed), coverage 85.07%, ruff/mypy clean. The two-server refusal is proven by a real integration test. Per the per-PR review norm, a quick focused review of the lock correctness + the CLI-mode implication.

362. user (2026-06-21T16:05:17.014Z)

One more consideration on the lock lifecycle: on the CLI single-shot path, main() returns normally (e.g. line 769, 837) without ever calling `release\_lock`. The lock relies on the OS releasing it at process exit. That's fine and documented ? but worth confirming there's no in-process workflow where main() returns and the process keeps running (e.g. main() called twice in one process, or a test importing main). The subprocess tests spawn fresh processes, so fine. In-process, `\_instance\_lock` would stay held until process exit; a second `main()` call in the same process would `acquire\_lock` the same path on a *new* handle ? which on POSIX (flock) is associated per-open-file-description so a new `open()` + flock would **block/fail** against the same process's existing lock (flock conflicts across different fds even in the same process), and on Windows msvcrt also conflicts across handles. This only matters for a hypothetical double-`main()`-in-one-process, not a real workflow. Noting as informational, not a finding.

I have enough to write the review.

Findings

[minor] backend/utils/single_instance.py:35 ? ``open(lock_path, "a+")`` is OUTSIDE the try/except. ****** If the lock directory/path is not writable (read-only app-home, permission error, path is a directory), ``open()`` raises ``OSError`/`PermissionError`` which propagates out of ``acquire_lock`` and crashes ``main()`` with an unhandled traceback ? instead of the clean ``[FATAL]`` ? Refusing to start message the design intends for the "can't run here" case. ``os.makedirs(..., exist_ok=True)`` at line 34 can raise too. ? ****Fix:**** wrap the ``makedirs`+`open`` in try/except and either return ``None`` (caller already prints FATAL + exits) or raise a typed error ``main()`` catches; at minimum document that ``acquire_lock`` can raise on an unwritable path so the caller isn't surprised.

[nit] backend/utils/single_instance.py:20,26 ? unparameterized ``IO`` generic ****** (``Optional[IO]``). mypy is reported clean (bare ``IO`` defaults to ``IO[Any]``), so this is cosmetic, but ``IO[str]`` is the precise type for an ``"a+"`/text-mode` handle. ? **\*\*Fix (optional):\*\*** annotate ``Optional[IO[str]]``.

[nit] backend/utils/single_instance.py:51 ? PID written without a trailing newline / encoding is platform-default. ****** ``fh.write(str(os.getpid()))`` is best-effort diagnostic only (the comment says so) and the test reads it back with ``strip()``, so harmless. Mentioned only for completeness; not worth changing.

Verified clean (no findings)

- ****Race prevention / ordering**** ? lock acquired at main.py:743, ``None``?``sys.exit(1)`` at 744-748, ****before**** ``Database()`` (750) and ``fail_stale_jobs()`` (754). The destructive sweep cannot run until the lock is held. Correct.
- ****fd kept alive**** ? ``_instance_lock`` is a module global (main.py:148), assigned under ``global _instance_lock`` (742); never GC'd ? lock held for process lifetime. Correct.
- ****OS release on crash/exit**** ? ``fcntl.flock`` (POSIX, line 43) and ``msvcrt.locking`` byte-range lock (Windows, line 40) are both kernel-managed and released when the handle closes or the process dies. The docstring's "no stale-lock problem" claim is accurate.
- ****CLI-mode implication**** ? lock is intentionally before the mode branch (738-748 precede the ``if args.video_path`` and not ``args.server`` at 759), so CLI single-shot ``*and*`` ``--server`` both lock. A CLI run while a server is up correctly refuses (it would otherwise sweep the server's jobs). No ``release_lock`` on the CLI return paths, but OS-on-exit covers it. ****No stale lock file on disk matters**** ? the lock is the open-handle, not file contents; the leftover ``lock`` file is harmless and re-lockable.
- ****Different output-dirs coexist**** ? lock path is ``db_path + ".lock"`` (743) and ``db_path`` = ``join(output_dir, db_name)`` (736); distinct output-dirs ? distinct lock paths ? no conflict. Covered by test_single_instance.py:30 and the integration test (test_server_entrypoint.py:136 uses same dir to prove refusal).
- ****msvcrt write-after-lock**** ? locks 1 byte at offset 0 (40), then ``seek(0)/truncate()/write`` (49-51) from the ``owning`` handle; Windows byte-range locks block ``other`` handles, not the owner, so the write succeeds. Correct.
- ****db_path is always local**** ? it is ``join(local output_dir, db_name)``; the ``gs://`/`gcs://`` URI handling applies only to the ``input video`` (``storage.resolve_input_ref``, 764-773), never to ``db_path``. The lock never sees a remote URI.
- ****Lock file ? SQLite DB file**** ? separate paths: ``sports_indexer.db`` vs ``sports_indexer.db.lock``. Holding the ``lock`` handle does not lock the DB file. Confirmed.
- ****sys.platform branching**** ? uses ``sys.platform == "win32"`` (correct for mypy platform-narrowing; the docstring at 22-23 explains why not ``os.name``); imports ``msvcrt`/`fcntl`` inside the matching branch only, so neither cross-errors. Correct.
- ****TOCTOU**** ? none: acquisition is a single atomic non-blocking lock syscall (``LOCK_NB`` / ``LK_NBLCK``); there is no check-then-act gap.

Verdict

Functionally correct and the race is genuinely closed; the only real (minor) issue is that an unwritable lock path raises out of ``acquire_lock`` instead of degrading to the clean FATAL refusal ? worth a try/except wrap, but not a merge blocker.

agentId: a0b5980ffee60fa70 (use SendMessage with to: 'a0b5980ffee60fa70' to continue this agent)

<usage>subagent_tokens: 50435

tool_uses: 10

duration_ms: 79823</usage>

363. assistant (2026-06-21T16:06:00.187Z)

Review verdict: race genuinely closed, all the cross-platform/CLI/lifecycle concerns clean. One minor fix worth folding ? if the lock path is unwritable, `acquire\_lock`'s `open()` raises and crashes `main()` instead of degrading. I'll make `main()` warn-and-proceed (a lock-infra failure shouldn't block a legitimate single instance, nor falsely claim "another instance"), plus the precise `IO[str]` annotation.

364. user (2026-06-21T16:06:05.303Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

365. user (2026-06-21T16:06:09.347Z)

The file C:\Users\avidu\Projects\bhi-wt-p0e\backend\utils\single_instance.py has been updated. All occurrences were successfully replaced. (file state is current in your context ? no need to Read it back)

366. user (2026-06-21T16:07:18.531Z)

```
..... [100%]
8 passed in 25.61s
=== ruff+mypy ===
All checks passed!
backend\pipeline\validators\base.py:36: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 2 source files
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

367. assistant (2026-06-21T16:07:22.929Z)

All green, mypy clean. Final full-suite confirmation, then commit P0e.

368. user (2026-06-21T16:09:00.150Z)

```
1277 passed, 7 skipped in 93.72s (0:01:33)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

369. assistant (2026-06-21T16:09:03.183Z)

Full suite green (1277 passed). Committing P0e and opening the PR.

370. user (2026-06-21T16:09:18.183Z)

```
=== status ===
M backend/main.py
M tests/test_server_entrypoint.py
?? backend/utils/single_instance.py
?? tests/test_single_instance.py
warning: in the working copy of 'backend/utils/single_instance.py', LF will be replaced by CRLF
the next time Git touches it
warning: in the working copy of 'tests/test_single_instance.py', LF will be replaced by CRLF the
next time Git touches it
=== HEAD ===
bd71048 feat(pkg): P0e ? single-instance lock so a 2nd process can't corrupt in-flight job state
(default-OFF)
branch 'feat/packaging-p0e-single-instance' set up to track 'origin/feat/packaging-p0e-single-
instance'.
```

To <https://github.com/Khelsutra/badminton-highlight-indexer.git>
* [new branch] feat/packaging-p0e-single-instance -> feat/packaging-p0e-single-instance
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

371. user (2026-06-21T16:09:34.700Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/298>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

372. assistant (2026-06-21T16:09:39.700Z)

**PR #298 opened: [feat(pkg): P0e ? single-instance
lock](<https://github.com/Khelsutra/badminton-highlight-indexer/pull/298>).** Watching CI.

373. user (2026-06-21T16:12:26.070Z)

Syntax sweep + import smoke (no GPU stack) pass 21s
<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379405>
Built wheel smoke (install + launch) pass 1m6s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379360>
Dependency license scan (no copyleft in runtime deps) pass 28s
<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379361>
Full test suite (offline, mocked) pass 2m37s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379376>
Golden fixtures cross-OS (macos-latest) pass 24s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379390>
Golden fixtures cross-OS (ubuntu-latest) pass 20s
<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379380>
Golden fixtures cross-OS (windows-latest) pass 1m18s
<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379367>
Mypy (lenient baseline) pass 37s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379454>
Ruff lint pass 8s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379382>
Syntax sweep + import smoke (no GPU stack) pass 21s
<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27910005616/job/82585379405>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

374. assistant (2026-06-21T16:12:30.055Z)

All 9 lanes green (incl. the built-wheel + license-scan lanes from P0d). Merging P0e ? the final P0 PR.

375. user (2026-06-21T16:12:39.753Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
e8bad40..c039b90 master -> origin/master
=== master ===
c039b90 feat(pkg): P0e ? single-instance lock so a 2nd process can't corrupt in-flight job state
(default-OFF) (#298)
=== PR #298 ===
MERGED
=== P0 PRs on master ===
c039b90 feat(pkg): P0e ? single-instance lock so a 2nd process can't corrupt in-flight job state
(default-OFF) (#298)
e8bad40 feat(pkg): P0d ? declare scipy (undeclared serving dep) + license-scan & built-wheel CI
lanes (#297)
```

```
3895963 feat(pkg): P0c ? config/profile-driven bind address + DNS-rebinding Host-header guard
(default-OFF) (#296)
fc78fce feat(pkg): P0b ? redact secret-shaped fields from the unauthenticated GET /api/config
(#295)
82a28f4 feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for config/.env/DB/output
(default-OFF) (#293)
```

376. assistant (2026-06-21T16:12:45.676Z)

? ****All 5 P0 PRs merged**** ? the P0 engine-prerequisite tranche is complete (master `c039b90`).
Let me finalize the tracker + memory.

377. user (2026-06-21T16:12:55.135Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

378. user (2026-06-21T16:13:10.957Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

379. assistant (2026-06-21T16:13:25.873Z)

Let me do a final state check.

380. user (2026-06-21T16:13:35.019Z)

```
=== open PRs (mine) ===
=== master HEAD ===
c039b90 feat(pkg): P0e ? single-instance lock so a 2nd process can't corrupt in-flight job state
(default-OFF) (#298)
=== my packaging worktrees cleaned? ===
(none ? all cleaned)
=== no GCP VMs (cost audit) ===
```

381. assistant (2026-06-21T16:13:55.936Z)

Clean checkpoint: ****0** open PRs, all worktrees cleaned, no GCP VMs running (\$0 spend ? P0 needed none).****** Here's where things landed.

? Plan delivered + all of P0 shipped

****The plan**** ? docs/PACKAGING_PLAN.md: one profile-selected build, cross-platform webserver (no native component), LIGHT-first with the full WASB/training/annotation path as the locked north-star (your 01/02 calls baked in). It's a local untracked tracked-doc + a memory checkpoint; say the word and I'll land it as a small docs PR too.

****P0 ? engine-prerequisite tranche, all 5 merged to master**** (`c039b90`), each isolated, default-OFF, adversarially reviewed, full-suite + 9-lane CI green, coverage held ~85%:

```
| PR | What it fixed | Proof |
|---|---|---|
| [#293 P0a](https://github.com/Khelsutra/badminton-highlight-indexer/pull/293) | **App-data HOME** (`RALLY_APP_HOME`) for config/.env/DB/output; fixed the doctor?server config-path disagreement (silent data-loss footgun) | e2e: state follows `RALLY_APP_HOME`, byte-identical when unset |
| [#295 P0b](https://github.com/Khelsutra/badminton-highlight-indexer/pull/295) | **Redacted `GET /api/config`** (was dumping the whole config unauthenticated) | full-`AppConfig` over-masking guard test |
| [#296 P0c](https://github.com/Khelsutra/badminton-highlight-indexer/pull/296) | **Config-
```

driven bind + DNS-rebinding Host guard** (default `127.0.0.1`) | middleware actually 400s a foreign Host |
| [#297 P0d](https://github.com/Khelsutra/badminton-highlight-indexer/pull/297) | **Declared `scipy`** (undeclared serving dep) + **license-scan & built-wheel CI lanes** | clean-venv wheel install + `indexer --help` |
| [#298 P0e](https://github.com/Khelsutra/badminton-highlight-indexer/pull/298) | **Single-instance OS lock** so a 2nd process can't corrupt in-flight jobs | two-server subprocess refusal test |

The engine is now **safe to launch from an arbitrary directory and to bundle** ? exactly what the distribution layer needs. CI gained two durable guardrails (a copyleft license scan and a real built-wheel smoke test).

Method notes (your working agreements)

- Worked each PR in an **isolated worktree off `origin/master`**, stage-explicit-paths, never touching the sibling sessions' WIP ? `origin/master` moved **6 times** under me during the session (43eccc1 ? c039b90); branching from the remote each time avoided any collision.
- **Adversarial review per PR** caught real issues I folded before merge (P0a's `load\_dotenv` byte-identity + `resolve\_output\_dir` normalization; P0c's edge-auth allowlist footgun ? promoted to a startup check; P0e's unwritable-path crash ? graceful degrade).

Next (your call)

P1 is the next tranche and the first genuinely cross-repo step: bundle the `khelsutra/web` SPA single-origin into the wheel + build the 3 net-new endpoints (`/api/capabilities`, `/api/videos` list, `/api/reels`) + a SPA?engine contract test. It touches the **separate `khelsutra` repo** (the SPA), so it needs a cross-repo build step and coordination ? a good place to confirm before I proceed. The GPU-validation work (P3/P4) is where I'll use your VM budget (confirm-cmd-and-cost, delete-after).

Want me to start **P1** now, or pause here?

382. user (2026-06-21T17:08:04.848Z)

I want you to continue but want to understand what do you mean by " It touches the separate khelsutra repo"

383. assistant (2026-06-21T17:08:32.662Z)

Good question ? let me ground the answer in the actual repo facts (and pull `khelsutra` first, since work now touches it).

384. user (2026-06-21T17:08:40.321Z)

=== khelsutra: pull + remote + head ===

Already up to date.

origin https://github.com/avidullu/khelsutra.git (fetch)

a2292ef feat(web): MULTILINGUAL Phase A ? Spanish/Danish/Indonesian SPA catalogs (#67)

=== where is the SPA? ===

NEXT_STEPS.md

README.md

deploy

docs

node_modules

site

web

--- web/ ---

README.md

dist

index.html

node_modules

package-lock.json

package.json

```

src
tsconfig.json
vite.config.ts

=== SPA build: outDir + scripts + API base ===
C:/Users/avidu/Projects/khelsutra/web/package.json:8:    "dev": "vite",
C:/Users/avidu/Projects/khelsutra/web/package.json:9:    "build": "vite build",
=== API base URL the SPA calls ===
1:// Typed client for the engine FastAPI. Default = SAME-ORIGIN relative `/api/*` (the single-
origin
2:// appliance reverse-proxies UI + engine behind one origin; in dev Vite proxies `/api` ?
:8000).
17:/** Resolve the engine API base. Default same-origin `"/api"`. For a SPLIT deploy ? SPA on
19: * set `VITE_API_BASE` at build time to the engine's API base **including the `/api` path**
20: * (e.g. `https://api.khelsutra.guru/api`), and lock the engine's `RALLY_CORS_ALLOW_ORIGINS`
to the
=== is web/dist committed or gitignored in khelsutra? ===
C:/Users/avidu/Projects/khelsutra/.gitignore:2:dist/
C:/Users/avidu/Projects/khelsutra/web/.gitignore:2:dist

=== engine: is backend/api/static the legacy/dead UI? ===
app.js
index.html
style.css

```

385. assistant (2026-06-21T17:10:03.049Z)

Pulled `khelsutra` (up to date, now at #67). Here's exactly what "touches the separate khelsutra repo" means, grounded in what I just verified:

The two-repo split

- **The engine** = **this** repo, `Khelsutra/badminton-highlight-indexer` ? the Python/FastAPI backend and the wheel I've been hardening.
- **The actual UI you click** (per-rally grid, query bar, build?download, en/kn/hi) = a **different git repo**, `avidullu/khelsutra`, in its `web/` folder (Vite/React/TS). Different repo ? different GitHub namespace even (`avidullu/` vs `Khelsutra/`), Cloudflare-connected, actively developed by the sibling sessions.
- What the engine serves today at `/` is a **legacy dead mini-dashboard** (`backend/api/static/{index.html,app.js,style.css}`), **not** that SPA.

Why P1 unavoidably reaches across repos

Right now those run as **two processes** in dev ? the Vite dev server proxies `/api` ? engine:8000. But a **batteries-included download** must be **one thing**: the engine serves the real UI itself, single-origin (the SPA already calls same-origin `/api`, so no SPA code change ? confirmed). To do that, the **engine wheel** has to physically contain the built UI.

The catch: the SPA's build output (`web/dist/`) is **gitignored in khelsutra** ? it's produced by `npm run build`, committed nowhere. And we can't run `npm` on the user's machine at `pip install` time. So **at our package-build time**, the SPA must already be built and copied into the engine tree (e.g. `backend/ui/`, force-included into the wheel). That build-time dependency on a second, independently-developed repo **is** the cross-repo coupling. It's not something the engine repo can decide entirely on its own ? and I don't want to unilaterally restructure your product/brand repo or its Cloudflare wiring.

So P1 splits cleanly:

- **P1a ? engine-only, fully unblocked**: the 3 net-new endpoints (`/api/capabilities`, `/api/videos` list, `/api/reels`) + an `/api/version` handshake + contract test. No khelsutra involvement. **I'll start here now.**
- **P1b ? the cross-repo bundling**: needs your call on **how** the SPA gets into the wheel.

That one decision is genuinely yours (it touches repo structure + your separate product repo), so:

386. user (2026-06-21T17:12:43.662Z)

Your questions have been answered: "How should the khelsutra SPA's built UI get into the engine wheel (so a downloaded engine serves the real UI single-origin)?"="1. with a goal to also be able to publish release (3) at some future point". You can now continue with these answers in mind.

387. assistant (2026-06-21T17:13:21.490Z)

Locked: ****option 1 (committed snapshot + refresh script) now, designed so it can also publish/consume a release artifact (option 3) later.**** I'll make the bundler source-agnostic from the start so that future swap is clean.

Continuing with ****P1a (engine-only, unblocked): the 3 endpoints + `/api/version`.** Setting up the worktree and grounding each endpoint in its real data source.

388. user (2026-06-21T17:13:26.977Z)

```
HEAD is now at c039b90 feat(pkg): P0e ? single-instance lock so a 2nd process can't corrupt in-
flight job state (default-OFF) (#298)
=== existing API endpoints ===
135:@app.get("/")
208:@app.get("/api/config"
218:@app.get("/api/health"
351:@app.post("/api/process"
402:@app.get("/api/jobs/{job_id}/cost"
417:@app.get("/api/videos/{video_id}/cost"
428:@app.get("/api/jobs/{job_id}"
461:@app.get("/api/highlights/{video_id}/{filename}"
485:@app.post("/api/query"
496:@app.post("/api/compile"

=== Database: video list + reels/highlights methods ===
339: def add_video(self, video_id: str, filepath: str, status: str, validation_results:
List[Dict[str, Any]]) -> bool:
358: def segment_watermarks(self, video_id: str) -> Tuple[int, int]:
372: def prune_segments(self, video_id: str, *, play_upto: Optional[int] = None,
393: def add_play_segment(self, video_id: str, start_time: float, end_time: float, server:
Optional[str] = None, receiver: Optional[str] = None, metadata: Optional[Dict[str, Any]] = None)
-> Optional[int]:
414: def add_candidate_segment(self, video_id: str, detector: str, start_time: float,
end_time: float,
445: def query_candidate_segments(self, video_id: str, detector: Optional[str] = None,
470: def get_video(self, video_id: str) -> Optional[Dict[str, Any]]:
479: def query_play_segments(self, video_id: str, filters: Dict[str, Any]) -> List[Dict[str,
Any]]:
529: def clear_video_data(self, video_id: str):
538: def create_job(self, job_id: str, video_path: str, segmenter: Optional[str] = None,
568: def set_job_video_id(self, job_id: str, video_id: str) -> bool:
677: def get_video_cost(self, video_id: str) -> Dict[str, Any]:

=== segmenter registry list + cuda probe + ffmpeg locate ===
C:/Users/avidu/Projects/bhi-wt-pla/backend/pipeline/segmenters/base.py:51: def
list_available(self) -> Dict[str, str]:
C:/Users/avidu/Projects/bhi-wt-pla/backend/config/presets.py:97: def probe_cuda_available() ->
bool:
```

389. assistant (2026-06-21T17:13:34.974Z)

No `list\_videos` or reels-list method exists yet. Let me read the videos table + `get\_video`, the highlights/compile handlers (reel storage), and the capability probes to ground the new endpoints.

390. user (2026-06-21T17:13:35.988Z)

```
470 def get_video(self, video_id: str) -> Optional[Dict[str, Any]]:
```



```

471         with self._connect() as conn:
472             row = conn.cursor().execute("SELECT * FROM videos WHERE id = ?",
(video_id,)).fetchone()
473             if row:
474                 d = dict(row)
475                 d["validation_results"] = json.loads(d["validation_results"] or "[]")
476                 return d
477             return None
478
479     def query_play_segments(self, video_id: str, filters: Dict[str, Any]) ->
List[Dict[str, Any]]:
480         """
481         Dynamically queries play segments based on filters:
482         - duration_min: float
483         - server: string
484         - receiver: string
485         - event_type: string (e.g. smash, foul)
486         """
487         query = "SELECT r.* FROM play_segments r WHERE r.video_id = ?"
488         params = [video_id]
489
490         if filters.get("duration_min"):
491             query += " AND r.duration >= ?"
492             params.append(filters["duration_min"])
493
494         if filters.get("server"):
495             query += " AND r.server = ?"
496             params.append(filters["server"])
497
498         if filters.get("receiver"):
499             query += " AND r.receiver = ?"

```

391. user (2026-06-21T17:13:36.885Z)

```

461     @app.get("/api/highlights/{video_id}/{filename}")
462     def download_highlight(video_id: str, filename: str):
463         """Stream a compiled highlight reel ? `filename` is the bare name given to
/api/compile
464         (which only writes into output_dir; the browser previously got back an unreachable
465         server-local path)."""
466         try:
467             name = safe_output_filename(filename)
468             except ValueError as e:
469                 raise HTTPException(status_code=400, detail=str(e)) from e
470             if not db_instance.get_video(video_id):
471                 raise HTTPException(status_code=404, detail="Indexed video record not found.")
472             output_dir = getattr(getattr(global_config, "output", None), "output_dir", "output")
or "output"
473             path = os.path.join(output_dir, name)
474             try:
475                 path = resolve_under_root(path, output_dir)
476             except ValueError as e:
477                 raise HTTPException(status_code=400, detail=str(e)) from e
478             if not os.path.isfile(path):
479                 raise HTTPException(status_code=404,
480                                     detail=f"No compiled file named '{name}'. Compile first via
/api/compile.")
481             media = "video/quicktime" if name.lower().endswith(".mov") else "video/mp4"
482             return FileResponse(path, media_type=media, filename=name)
483
484
485     @app.post("/api/query")
486     def query_play_segments(req: QueryRequest):
487         filters = {
488             "server": req.server,

```

```

489         "receiver": req.receiver,
490         "duration_min": req.duration_min,
491         "event_type": req.event_type
492     }
493     segments = db_instance.query_play_segments(req.video_id, filters)
494     return {"play_segments": segments}
495
496 @app.post("/api/compile")
497 def compile_highlights(req: CompileRequest):
498     video = db_instance.get_video(req.video_id)
499     if not video:
500         raise HTTPException(status_code=404, detail="Indexed video record not found.")
501
502     # Retrieve matching play segments from db
503     all_segments = db_instance.query_play_segments(req.video_id, {})
504     matched_segments = [s for s in all_segments if s["id"] in req.segment_ids]
505

```

392. assistant (2026-06-21T17:13:45.210Z)

Let me read the full compile handler to see exactly where reels land and whether they're associated with a `video\_id` (critical for designing `/api/reels/{video\_id}`).

393. user (2026-06-21T17:13:45.821Z)

```

496     @app.post("/api/compile")
497     def compile_highlights(req: CompileRequest):
498         video = db_instance.get_video(req.video_id)
499         if not video:
500             raise HTTPException(status_code=404, detail="Indexed video record not found.")
501
502         # Retrieve matching play segments from db
503         all_segments = db_instance.query_play_segments(req.video_id, {})
504         matched_segments = [s for s in all_segments if s["id"] in req.segment_ids]
505
506         if not matched_segments:
507             raise HTTPException(status_code=400, detail="No matching play segments found to
508 stitch.")
509
510         # Reject path traversal / absolute output names (os.path.join would otherwise let an
511         # absolute or '..'-laden output_filename write anywhere on disk).
512         try:
513             out_name = safe_output_filename(req.output_filename)
514         except ValueError as e:
515             raise HTTPException(status_code=400, detail=str(e)) from e
516
517         # The configured shot-count floor (stitching.min_shot_count) applies on the API
518         # path too ? parity with scripts/run_process_and_stitch.py. Segments WITHOUT
519         shot_count
520         # metadata are exempt: the floor filters known counts only, so explicitly
521         # requested manual/legacy segments can't silently vanish under the default floor.
522         stitching = getattr(global_config, "stitching", None) or StitchingConfig()
523         kept = []
524         for s in matched_segments:
525             meta = s.get("metadata") or {}
526             sc = meta.get("shot_count") if isinstance(meta, dict) else None
527             try:
528                 if sc is not None and int(sc) < stitching.min_shot_count:
529                     continue
530             except (TypeError, ValueError):
531                 pass # unparseable count -> treat as unknown, keep
532             kept.append(s)
533         if len(kept) < len(matched_segments):
534             logger.info(f"[compile] stitching.min_shot_count={stitching.min_shot_count}:
535 dropped "

```

```

533             f"{len(matched_segments) - len(kept)}/{len(matched_segments)}
segments below the floor.")
534         if not kept:
535             raise HTTPException(status_code=400,
536                                 detail=f"All {len(matched_segments)} matching segments fall
below "
537 f"stitching.min_shot_count={stitching.min_shot_count}.")
538
539         # Extract start/end time pairs
540         time_pairs = [(s["start_time"], s["end_time"]) for s in kept]
541
542         # Run compiler with the configured stitching paddings + optional branding watermark
543         output_dir = getattr(getattr(global_config, "output", None), "output_dir", "output")
or "output"
544         compiler = VideoCompiler(output_dir, begin_padding=stitching.begin_padding,
545                                 end_padding=stitching.end_padding,
watermark=stitching.watermark)
546         try:
547             # The stored `filepath` is the SOURCE ref ? a bucket/Drive URI (gs://, s3://,
gdrive://) for a
548             # video ingested via /api/process from a URI, or a plain local path. Resolve it
to a LOCAL file
549             # the stitcher can read: identity for a local path (byte-identical to before),
fetch-to-scratch
550             # for a bucket URI (the same seam the process path uses, _execute_processing
main.py:159).
551             # Without this, compiling a bucket-ingested video 500s ? ffmpeg can't open a
gs:// path.
552             local_src = storage.acquire_local_input(video["filepath"],
getattr(global_config, "storage", None))
553             out_path = compiler.compile_highlights(local_src, time_pairs, out_name)
554             # `filepath` is server-local (not reachable by the browser); also return the
ready-to-use
555             # download URL so the SPA needn't hand-build it (P6 / PER_RALLY_SELECTION.md
$4.3).

```

394. assistant (2026-06-21T17:15:03.829Z)

Important finding (verify-before-scoping): **reels are stored flat in `output\_dir/<filename>` with no `video\_id` association** ? `/api/compile` writes a bare filename and `/api/highlights` ignores the `video\_id` for the path. So a per-video `/api/reels/{video\_id}` would be a phantom. I'll design `/api/reels` honestly: a global list of compiled reels on disk + a `video\_id`-free download, matching the actual flat storage. Let me check the videos-table columns + main.py imports to finish the design.

395. user (2026-06-21T17:15:07.097Z)

```

=== videos table columns (migrations) ===
118:         CREATE TABLE IF NOT EXISTS videos (
232:         self._add_column_idempotent(cursor, "ALTER TABLE videos ADD COLUMN user_id TEXT")
350:         "INSERT INTO videos (id, filepath, status, validation_results) VALUES (?, ?, ?,
?) "
=== add_video INSERT columns ===
def add_video(self, video_id: str, filepath: str, status: str, validation_results:
List[Dict[str, Any]]) -> bool:
    """Register or update a video row WITHOUT touching its child segments.

    Uses UPSERT, not INSERT OR REPLACE: REPLACE deletes the conflicting row, which ?
with foreign_keys ON ? cascade-deletes every play_segment/event/candidate_segment.
That made a status update (e.g. marking 'failed' on re-validation, or 'processed'
before the segmenter runs) silently destroy prior good results. UPSERT mutates the
row in place; clearing segments is now an explicit, deliberate operation
(clear_video_data / prune_segments).
    """

```

```

return self._execute_write(
    "INSERT INTO videos (id, filepath, status, validation_results) VALUES (?, ?, ?, ?) "
    "ON CONFLICT(id) DO UPDATE SET "
    "filepath=excluded.filepath, status=excluded.status, "
    "validation_results=excluded.validation_results",
    (video_id, filepath, status, json.dumps(validation_results)),
    "Failed to add video",
)

```

=== main.py: shutil / importlib.metadata / probe_cuda / segmenter_registry imports ===

```

35:from backend.pipeline.segmenters.base import segmenter_registry
283:    segmenter_cls = segmenter_registry.get(seg_name)
287:    available = list(segmenter_registry.list_available().keys())
384:    if req.segmenter and not segmenter_registry.get(req.segmenter):
385:        available = list(segmenter_registry.list_available().keys())
662:    available_segmenters = list(segmenter_registry.list_available().keys())
813:    segmenter_cls = segmenter_registry.get(seg_name)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

396. assistant (2026-06-21T17:15:12.590Z)

Design complete. Videos table = `id, filepath, status, validation\_results, user\_id` (no timestamp ? order by rowid). Adding `list\_videos()` to the DB.

397. user (2026-06-21T17:15:19.060Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\infrastructure\database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

398. assistant (2026-06-21T17:15:23.158Z)

Now the main.py imports + the 4 endpoints.

399. user (2026-06-21T17:15:25.120Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

400. user (2026-06-21T17:15:31.451Z)

```

34    from fastapi.responses import FileResponse, HTMLResponse, JSONResponse
35    from backend.pipeline.validators.base import registry
36    from backend.infrastructure.database import Database
37    from backend.pipeline.segmenters.base import segmenter_registry
38    from backend.pipeline.stitcher.compiler import VideoCompiler
39    from backend.utils.hashing import compute_video_id # canonical file-identity hashing
40    from backend.utils.paths import resolve_under_root, safe_output_filename,
validate_input_path
41    from backend.utils.redact import redact_secrets # P0b: mask secret-shaped fields in
/api/config
42    from backend.utils.single_instance import acquire_lock # P0e: one instance per database
43    from backend.config import ( # new typed config
44        load_config as load_app_config,
45        AppConfig,

```

401. user (2026-06-21T17:15:34.874Z)

```

43    from backend.config import ( # new typed config
44        load_config as load_app_config,
45        AppConfig,
46        StitchingConfig,
47        enforce_startup_checks,
48        StartupCheckError,
49    )

```

```
50         from backend.results import Failure # standardized error/result contract
```

402. user (2026-06-21T17:15:39.923Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

403. assistant (2026-06-21T17:15:47.164Z)

Now add `API\_VERSION` + the four endpoints. Let me find the anchor after `/api/health`.

404. user (2026-06-21T17:15:47.583Z)

```
219         return redact_secrets(global_config.model_dump(mode="json"))
220
221     @app.get("/api/health")
222     def health():
223         """Cheap liveness/readiness probe (no auth) for the SPA's offline banner
224         (HOSTING_PLAN.md:42). Reports the configured sport + default segmenter so the UI can
225         confirm the engine is reachable and which detector is wired. Never raises."""
226         sport = getattr(global_config, "sport", None)
227         indexing = getattr(global_config, "indexing", None)
228         return {
229             "status": "ok",
230             "sport": sport,
231             "default_segmenter": getattr(indexing, "default_segmenter", None),
232         }
233
234     def _execute_processing(req: ProcessRequest, progress=None) -> Dict[str, Any]:
```

405. user (2026-06-21T17:16:12.978Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

406. assistant (2026-06-21T17:16:37.959Z)

Now the tests. Let me find the anchor in `test\_api\_endpoints.py` (the config tests merged from P0b).

407. user (2026-06-21T17:16:39.607Z)

```
28: def env(tmp_path, monkeypatch):
61: def test_api_config_ok(env):
68: def test_api_config_redacts_secret_shaped_fields(env):
81: def test_process_success_with_mocked_segmenter(env):
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

408. user (2026-06-21T17:16:43.833Z)

```
68     def test_api_config_redacts_secret_shaped_fields(env):
69         """P0b: the unauthenticated /api/config must never serialize a credential-shaped
value.
70         AppConfig stores no secrets today, so inject one via its extra='allow' tolerance and
assert
71         it is masked while non-secret fields are served unchanged."""
72         client, db, tmp_path, mp = env
73         mp.setattr(m, "global_config", AppConfig(gemini_api_key="SECRET-LEAK-123"))
74         r = client.get("/api/config")
75         assert r.status_code == 200
76         assert "SECRET-LEAK-123" not in r.text # the secret never crosses the wire
77         assert r.json()["gemini_api_key"] == "****REDACTED****"
78         assert r.json()["sport"] == "badminton" # non-secret fields still served
79
```

409. user (2026-06-21T17:17:01.375Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\tests\test_api_endpoints.py has been updated successfully. (file state is current in your context ? no need to Read it back)

410. user (2026-06-21T17:17:22.966Z)

```
..... [100%]
38 passed in 13.84s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

411. assistant (2026-06-21T17:17:29.088Z)

All 38 pass. Now lint/type + full suite + coverage.

412. user (2026-06-21T17:20:38.418Z)

```
=== RUFF ===
All checks passed!
=== MYPY ===
backend\pipeline\validators\base.py:36: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 2 source files
=== FULL SUITE + COV ===
backend\infrastructure\database.py          355      30      92%
backend\main.py                             519     143      72%
TOTAL                                     9826    1468      85%
Required test coverage of 70% reached. Total coverage: 85.06%
1284 passed, 7 skipped, 7 warnings in 146.55s (0:02:26)
  RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

413. assistant (2026-06-21T17:20:49.179Z)

Full suite green (1284 passed), ruff/mypy clean, coverage 85.06%, `database.py` 92%. Focused adversarial review of the 4 new endpoints before commit.

414. user (2026-06-21T17:24:35.424Z)

Path safety holds empirically:

- `..%2f..%2f...` (encoded `/`) ? Starlette decodes `%2f` to `/` inside the single-segment param, then `safe\_output\_filename` would reject it ? but here it 404s (the decoded `../../backend/main.py` doesn't match a `.mp4`/`.mov` and isn't a file in output_dir; the path containment + ext gate catch it). Either way, no escape.
- Null byte ? 400 (rejected by `safe\_output\_filename`).
- Backslash ? 400 (rejected).

Traversal is not reachable. Q2 is clean.

I have enough evidence for all six review axes. Here are my findings.

Findings

****[minor] backend/main.py:49 and backend/main.py:261 ?** The inline comments `# cheap env-based GPU probe (no torch import)` and `cuda = probe\_cuda\_available() # cheap, env-based (no torch import)` are factually false. **** `probe\_cuda\_available()` ?** `presets.py:106 \_cuda\_available()` ? `import torch; torch.cuda.is\_available()` (presets.py:90-92). Measured: first

`/api/capabilities` call takes ~3.5s (torch import + CUDA init); subsequent calls are cached (0ms). The function is correctly wrapped in try/except so it never raises and returns the right bool, but it is NOT cheap and NOT env-based on the first call. ****Fix:**** correct both comments to "best-effort CUDA probe (lazy torch import, ~3s first call, cached after; never raises)", or, if first-run-wizard latency matters, gate the torch probe behind a query param / cache it at startup.

****[minor]** backend/infrastructure/database.py:480 ? `list\_videos` docstring claims "the videos table has no timestamp, so order by rowid", but the table DOES have `processed\_at` TIMESTAMP DEFAULT CURRENT_TIMESTAMP` (database.py:121). **\*\*** The ordering choice itself is *\*correct and arguably better\** ? `id` is TEXT PRIMARY KEY so `rowid` is a true monotonic insertion counter, immune to the 1-second granularity ties `processed_at` would suffer. Only the stated rationale is wrong. ****Fix:**** reword to "order by rowid (monotonic insertion order; more reliable than the 1s-granularity processed_at for newest-first)".

****[nit]** backend/main.py:269 ? `/api/capabilities` serializes full `ffmpeg\_path` / `ffprobe\_path` on an unauthenticated endpoint, disclosing the install layout (and OS username on Windows paths). ****** This is **consistent** with the established posture of the sibling localhost endpoints ? `/api/videos` returns `filepath` (full source paths) and `/api/config` returns `output.output\_dir` ? so it is an acceptable disclosure under the current single-user-localhost threat model, not a regression. No secret value or full secret path is leaked: `gemini\_key\_present` is `bool(os.environ.get("GEMINI\_API\_KEY"))` and `wasb\_weights\_present` is `bool(weights\_path)` and `os.path.exists(weights\_path)` ? both pure bools, neither the key nor the weights path crosses the wire (verified by the test asserting `"AIza-should-not-leak" not in r.text`). ****Fix (optional):**** if exposure tightens later, drop `ffmpeg\_path`/`ffprobe\_path` to bools-only like the secret fields; fine to defer.

Axes that are clean (no issue)

- ****Q1 SECRET LEAK:**** Clean. Only booleans for key/weights presence; no key value, no full secret path. (`ffmpeg\_path` disclosure noted above as a consistent nit.)
- ****Q2 PATH SAFETY:**** Clean. `safe\_output\_filename` rejects empty/null-byte/`/`/`\\`/absolute/`../`/non-basename; `resolve\_under\_root` (realpath + commonpath) is real defense-in-depth catching symlink escape; the `.mp4`/`.mov` ext gate is ANDed with `os.path.isfile`. Route is single-segment `{filename}` (not `{filename:path}`) so separators don't span. Empirically probed encoded-`/`, encoded-`..`, null-byte, backslash ? all 400/404, no escape. `os.scandir` handle is closed via `with entries:` (main.py:316). No exploitable TOCTOU (realpath resolved, then `isfile`, then `FileResponse` opens ? worst case is a 404/serve of a since-replaced file under the same contained root, not an escape).
- ****Q3 HONESTY:**** Accurate. Confirmed `/api/compile` ? `VideoCompiler(output\_dir).compile\_highlights` writes flat to `os.path.join(self.output\_dir, os.path.basename(name))` (compiler.py:326) with no `video\_id` association ? the "global flat list / per-video would be a phantom" docstring is honest. `e.is\_file()` correctly skips the per-video `tempfile.TemporaryDirectory(dir=output\_dir)` scratch subdirs (compiler.py:340), and the SQLite DB / `.lock` file are excluded by the `.mp4`/`.mov` ext gate regardless of location. Test covers `notes.txt` + a subdir exclusion.
- ****Q4 /api/videos:**** Clean. SQL parameterized ? `LIMIT` ? with `int(limit)` coercion (database.py:485-487), no injection. `SELECT \*` + `validation\_results` JSON-parse mirrors `get\_video` exactly, so it leaks nothing the existing per-row endpoint didn't already expose. `ORDER BY rowid DESC` is sound (see minor above).
- ****Q5 ROBUSTNESS:**** Holds the "never raises" contract. All config access is None-safe `getattr` chains (verified `global\_config=None` path), `segmenter\_registry.list\_available()` is per-item try/excepted and returns `{}` when empty, `probe\_cuda\_available()` is try/excepted, `shutil.which`/`os.path.exists("")` don't raise. (The torch-import **cost**, not correctness, is the only issue ? see minor.)
- ****Q6 /api/version:**** Correct. `\_engine\_version()` catches `PackageNotFoundError` ? `"unknown"` for editable/source runs (main.py:240-244). `API\_VERSION = 1` contract comment ("BUMP when an existing endpoint's shape changes incompatibly; additive endpoints need no bump") is accurate and matches the additive nature of this PR.

Verdict

Clean ? no blockers or majors; two minor doc/comment-accuracy fixes (the false "no torch import" comments at main.py:49/261 are the most worth correcting since they misrepresent a real ~3.5s first-call cost) and one optional nit. Safe to merge after the comment fix.

agentId: a77c136ef76cc9f16 (use SendMessage with to: 'a77c136ef76cc9f16' to continue this agent)
<usage>subagent_tokens: 69075

tool_uses: 27
duration_ms: 213284</usage>

415. assistant (2026-06-21T17:25:08.707Z)

Two fair minor fixes from the review: (1) my "cheap, no-torch" comment on `probe\_cuda\_available()` is **wrong** ? it lazily imports torch (~3.5s first call) and, worse, would false-negative in the LIGHT tier where torch is absent. The honest cheap signal for the wizard is `nvidia-smi` presence (what `doctor.py` uses). (2) the `list\_videos` docstring claims "no timestamp" but the table has `processed\_at`. Fixing both.

416. user (2026-06-21T17:25:10.763Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

417. user (2026-06-21T17:25:24.945Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

418. user (2026-06-21T17:25:29.594Z)

```
<tool_use_error>String to replace not found in file.  
String:      assert isinstance(body["gpu"]["cuda_available"], bool)</tool_use_error>
```

419. assistant (2026-06-21T17:25:33.633Z)

That assertion is in the test file. Fixing it there + the docstring.

420. user (2026-06-21T17:25:35.024Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\tests\test_api_endpoints.py has been updated successfully. (file state is current in your context ? no need to Read it back)

421. user (2026-06-21T17:25:41.469Z)

The file C:\Users\avidu\Projects\bhi-wt-p1a\backend\infrastructure\database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

422. user (2026-06-21T17:26:12.067Z)

```
..... [100%]  
38 passed in 12.31s  
=== ruff+mypy ===  
All checks passed!  
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not  
checked, consider using --check-untyped-defs [annotation-unchecked]  
Success: no issues found in 1 source file  
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

423. assistant (2026-06-21T17:26:17.920Z)

Review fixes folded, all green. Final full-suite confirmation, then commit P1a.

424. user (2026-06-21T17:28:00.134Z)

```
1284 passed, 7 skipped in 98.62s (0:01:38)  
=====
```

RESULT: PASS - all tests green

```
=====
```

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

425. assistant (2026-06-21T17:28:03.193Z)

Full suite green (1284). Committing P1a and opening the PR.

426. user (2026-06-21T17:28:18.536Z)

```
=== status ===
M backend/infrastructure/database.py
M backend/main.py
M tests/test_api_endpoints.py
=== HEAD ===
3524ab5 feat(pkg): P1a ? net-new console/wizard endpoints: /api/version, /api/capabilities,
/api/videos, /api/reels
branch 'feat/packaging-p1a-endpoints' set up to track 'origin/feat/packaging-p1a-endpoints'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch]      feat/packaging-p1a-endpoints -> feat/packaging-p1a-endpoints
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

427. user (2026-06-21T17:28:36.403Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/300>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

428. assistant (2026-06-21T17:28:40.237Z)

PR #300 opened: [feat(pkg): P1a ? console/wizard endpoints](<https://github.com/Khelsutra/badminton-highlight-indexer/pull/300>). Watching CI.

429. user (2026-06-21T17:31:27.012Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    20s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689286
Built wheel smoke (install + launch)            pass    1m9s   https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689263
Dependency license scan (no copyleft in runtime deps) pass    30s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689450
Full test suite (offline, mocked)                pass    2m26s   https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689244
Golden fixtures cross-OS (macos-latest)          pass    25s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689254
Golden fixtures cross-OS (ubuntu-latest)         pass    33s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689253
Golden fixtures cross-OS (windows-latest)        pass    1m15s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689266
Mypy (lenient baseline) pass    28s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689249
Ruff lint    pass    11s    https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689335
Syntax sweep + import smoke (no GPU stack)      pass    20s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912029819/job/82590689286
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

430. assistant (2026-06-21T17:31:33.386Z)

All 9 lanes green. Merging P1a and cleaning up.

431. user (2026-06-21T17:31:40.800Z)

failed to run git: fatal: 'master' is already used by worktree at

'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
cb7394b..be0c434 master -> origin/master
=== master ===
be0c434 feat(pkg): P1a ? net-new console/wizard endpoints: /api/version, /api/capabilities,
/api/videos, /api/reels (#300)
=== PR #300 ===
MERGED
```

432. assistant (2026-06-21T17:32:09.423Z)

P1a ****merged**** (master `be0c434`). Updating tracker + memory, then starting ****P1b** (the cross-repo SPA bundling)******.

433. user (2026-06-21T17:32:17.531Z)

```
<tool_use_error>String to replace not found in file.
String: | **P1a** | **Bundle the SPA single-origin**: build `khelsutra/web` at package time ?
vendor `web/dist/` into the wheel (hatchling force-include -> `backend/ui/`) + SPA history-
fallback; retire the legacy static mount; record bundle hash | P0a | No | ? | ? |
| P1b | Build `GET /api/capabilities` (honest hardware/segmenter/ffmpeg/weights probe), `GET
/api/videos` (list; today only single-video get), `GET /api/reels/{video_id}` (reload-survival)
| P1a | No | ? | ? |
(note: Edit also tried swapping \uxxxx escapes and their characters; neither form matched, so
the mismatch is likely elsewhere in old_string. Re-read the file and copy the exact surrounding
text.)</tool_use_error>
```

434. assistant (2026-06-21T17:32:23.415Z)

Let me read the actual P1 rows in the doc.

435. user (2026-06-21T17:32:24.463Z)

```
96:| **P1a** | **Bundle the SPA single-origin**: build `khelsutra/web` at package time ? vendor
`web/dist/` into the wheel (hatchling force-include ? `backend/ui/`) + SPA history-fallback;
retire legacy static mount; record bundle hash | P0a | No | ? | ? |
97:| P1b | **Net-new endpoints** the wizard/console need: `GET /api/capabilities` (honest
hw/segmenter/ffmpeg/weights probe), `GET /api/videos` (list), `GET /api/reels/{video_id}` +
`/api/version` handshake + SPA?engine contract test | P1a | No | ? | ? |
99:| P2b | **Cross-platform launcher + install path** (Win/Mac/Linux, no native component): `uv
tool install` / generated script + a pure-Python `indexer --app` (health-poll ?
`webbrowser.open`; free-port fallback; graceful shutdown); per-OS ffmpeg acquisition (winget /
brew / apt) via the P2a resolver; ship `config.json` defaulting truly-local; optional thin per-
OS shortcut | P2a, P1a | No (code-signing only if clickable binaries are ever shipped ? 09) | ?
| ? |
100:| P2c | **In-UI 3-step first-run wizard** (Gemini key / Compute / Videos), rendered from
`/api/capabilities`, en/kn/hi | P1b | No | ? | ? |
104:| P3c | **"Push to cloud" toggle** (flip cloud-serving + confirm + INR estimate + per-video
cost cap) + wire worker usage-emission into the M8 ledger | P1b | Owner (cloud spend) | ? | ? |
124:- [ ] P1a/P1b merged: FastAPI serves the **real SPA** single-origin with refresh-survival;
capabilities/videos/reels return real data; contract test guards SPA?engine skew.
```

436. user (2026-06-21T17:32:36.566Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

437. user (2026-06-21T17:32:52.348Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

438. assistant (2026-06-21T17:33:02.475Z)

Now scout the SPA build (measure the snapshot size + pin the commit) before committing assets.
Building `khelsutra/web` fresh.

439. user (2026-06-21T17:33:07.747Z)

```
=== khelsutra commit to pin ===
46e0244b0c862a4a9d845b8947843b5dfa1f2d9f
46e0244 feat(web): A9 ? localized tap-only filter controls (#68)
=== node + build the SPA ===
v26.1.0
rendering chunks...
computing gzip size...
dist/index.html
  0.40 kB ? gzip:  0.27 kB
dist/assets/noto-sans-devanagari-latin-ext-wght-normal-CLKupaxV.woff2
  13.85 kB
dist/assets/noto-sans-kannada-latin-ext-wght-normal-MdBBF41x.woff2
  13.99 kB
dist/assets/noto-sans-devanagari-latin-wght-normal-Bxfr4UJH.woff2
  25.38 kB
dist/assets/noto-sans-kannada-latin-wght-normal-DeLZCgFT.woff2
  27.58 kB
dist/assets/noto-sans-kannada-kannada-wght-normal-DuwB-P8d.woff2
  90.80 kB
dist/assets/noto-sans-devanagari-devanagari-wght-normal-DWI_L5Cu.woff2
 121.19 kB
dist/assets/index-FJh_kElx.css
  24.83 kB ? gzip:  5.43 kB
dist/assets/index-ZjpoX6Fg.js
 373.80 kB ? gzip: 117.33 kB
? built in 1.26s
```

440. assistant (2026-06-21T17:34:12.045Z)

SPA builds cleanly ? snapshot is ****~692 KB total**** (index.html + 373 KB JS + 25 KB CSS + ~293 KB of Noto woff2 fonts). That's small enough to commit comfortably (option 1). Pinning khelsutra `46e0244`. Let me measure exactly + set up the engine worktree.

441. user (2026-06-21T17:34:16.587Z)

```
=== dist size + file list ===
701K   C:/Users/avidu/Projects/khelsutra/web/dist
assets/index-FJh_kElx.css
assets/index-ZjpoX6Fg.js
assets/noto-sans-devanagari-devanagari-wght-normal-DWI_L5Cu.woff2
assets/noto-sans-devanagari-latin-ext-wght-normal-CLKupaxV.woff2
assets/noto-sans-devanagari-latin-wght-normal-Bxfr4UJH.woff2
assets/noto-sans-kannada-kannada-wght-normal-DuwB-P8d.woff2
assets/noto-sans-kannada-latin-ext-wght-normal-MdBBF41x.woff2
assets/noto-sans-kannada-latin-wght-normal-DeLZCgFT.woff2
index.html
```

```
=== engine API_VERSION (for the contract) ===
```

```
=== create P1b worktree ===
```

```
HEAD is now at be0c434 feat(pkg): P1a ? net-new console/wizard endpoints: /api/version,
/api/capabilities, /api/videos, /api/reels (#300)
be0c434 feat(pkg): P1a ? net-new console/wizard endpoints: /api/version, /api/capabilities,
/api/videos, /api/reels (#300)
```

442. assistant (2026-06-21T17:34:51.545Z)

dist = 701 K, 9 files ? comfortable to commit. Worktree ready off master (which has P1a's `API\_VERSION`). Writing the source-agnostic bundler (build-from-checkout now; tarball/release-

artifact path stubbed for option 3 later).

443. assistant (2026-06-21T17:35:12.384Z)

```
#!/usr/bin/env bash
```

Bundle the KhelSutra web SPA into the engine so a downloaded engine serves the ready-to-run single-origin (PACKAGING_PLAN.md P1b).

```
#
```

OPTION 1 (now, owner-chosen): build the SPA from a khelsutra checkout and vendor it into `backend/ui/`, which is COMMITTED ? so the wheel + CI build fully offline/reproducible the snapshot by re-running this on an SPA release; commit the diff.

```
#
```

OPTION 3 (designed-for, future): point --from at a published dist tarball/URL (a GitHub artifact) to vendor a pre-built SPA without a local khelsutra checkout or Node. Already

```
#
```

Usage:

scripts/bundle_ui.sh [SOURCE]

SOURCE = a khelsutra checkout dir (default: \$KHELSUTRA_DIR or ../khelsutra) [builds SPA] | a path/URL ending .tgz/.tar.gz [unpacks it]

BUILD=0 scripts/bundle_ui.sh <khelsutra-dir> # skip npm; reuse an existing web/dist

```
set -euo pipefail
```

```
ENGINE_ROOT="$(cd "$(dirname "${BASH_SOURCE[0]}")/.." && pwd)"
```

```
UI_DEST="$ENGINE_ROOT/backend/ui"
```

```
SOURCE="${1:-${KHELSUTRA_DIR:-$ENGINE_ROOT/../khelsutra}}"
```

```
stage_dist() { # populates $1 (a dist dir) into a clean $UI_DEST
```

```
    local dist="$1"
```

```
    [ -f "$dist/index.html" ] || { echo "ERROR: no index.html in $dist ? not a built SPA dist"
```

```
>&2; exit 1; }
```

```
    rm -rf "$UI_DEST"
```

```
    mkdir -p "$UI_DEST"
```

```
    cp -r "$dist"/. "$UI_DEST"/
```

```
}
```

```
SPA_REPO="unknown"
```

```
SPA_COMMIT="unknown"
```

```
case "$SOURCE" in
```

```
    *.tgz|*.tar.gz|http:/*|https:/*)
```

```
        # OPTION 3: a published dist artifact.
```

```
        tmp="$(mktemp -d)"; trap 'rm -rf "$tmp"' EXIT
```

```
        if [[ "$SOURCE" == http:/*/* ]]; then
```

```
            echo "[bundle_ui] fetching dist artifact: $SOURCE"
```

```
            curl -fsSL "$SOURCE" -o "$tmp/dist.tgz"; SRC_TGZ="$tmp/dist.tgz"; SPA_REPO="$SOURCE"
```

```
        else
```

```
            SRC_TGZ="$SOURCE"; SPA_REPO="$SOURCE"
```

```
        fi
```

```
        tar -xzf "$SRC_TGZ" -C "$tmp"
```

```
        # accept either dist/* at the root or a top-level dir containing index.html
```

```
        if [ -f "$tmp/index.html" ]; then DIST="$tmp"; else DIST="$(dirname "$(find "$tmp" -name index.html | head -1)");" ; fi
```

```
        stage_dist "$DIST"
```

```
        ;;
```

```
    *)
```

```
        # OPTION 1: a khelsutra checkout ? build it.
```

```

WEB="$SOURCE/web"
[ -d "$WEB" ] || { echo "ERROR: $WEB not found (pass a khelsutra checkout dir or a dist
tarball)" >&2; exit 1; }
SPA_REPO="$(git -C "$SOURCE" config --get remote.origin.url 2>/dev/null || echo "$SOURCE")"
SPA_COMMIT="$(git -C "$SOURCE" rev-parse HEAD 2>/dev/null || echo unknown)"
if [ "${BUILD:-1}" = "1" ]; then
    echo "[bundle_ui] building SPA in $WEB (npm ci && npm run build)"
    ( cd "$WEB" && npm ci && npm run build )
else
    echo "[bundle_ui] BUILD=0 ? reusing existing $WEB/dist"
fi
stage_dist "$WEB/dist"
;;
esac

```

Provenance + the SPA<->engine contract version (the engine's current API_VERSION

```

ENGINE_API_VERSION="$(grep -E '^API_VERSION\s*=' "$ENGINE_ROOT/backend/main.py" | grep -oE
'[0-9]+' | head -1)"
CONTENT_HASH="$(find "$UI_DEST" -type f ! -name ui_version.json | LC_ALL=C sort | xargs
sha256sum | sha256sum | cut -d ' ' -f1)"
BUILT_AT="$(date -u +%Y-%m-%dT%H:%M:%SZ 2>/dev/null || echo unknown)"

```

```

cat > "$UI_DEST/ui_version.json" <<JSON
{
  "spa_repo": "$SPA_REPO",
  "spa_commit": "$SPA_COMMIT",
  "engine_api_version": ${ENGINE_API_VERSION:-null},
  "content_hash": "$CONTENT_HASH",
  "built_at": "$BUILT_AT"
}
JSON

```

```

echo "[bundle_ui] vendored SPA -> $UI_DEST (spa_commit=$SPA_COMMIT,
engine_api_version=$ENGINE_API_VERSION)"
echo "[bundle_ui] files:"; find "$UI_DEST" -type f | sed "s|$UI_DEST/||" | sort

```

444. user (2026-06-21T17:35:12.426Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p1b\scripts\bundle_ui.sh (file state is current in your context ? no need to Read it back)

445. assistant (2026-06-21T17:35:16.551Z)

Now run the bundler (reusing the dist I just built) to produce the committed `backend/ui/` snapshot.

446. user (2026-06-21T17:35:18.863Z)

```

[bundle_ui] BUILD=0 ? reusing existing C:/Users/avidu/Projects/khelsutra/web/dist
[bundle_ui] vendored SPA -> /c/Users/avidu/Projects/bhi-wt-p1b/backend/ui
(spa_commit=46e0244b0c862a4a9d845b8947843b5dfa1f2d9f, engine_api_version=1)
[bundle_ui] files:
assets/index-FJh_kElx.css
assets/index-ZjpoX6Fg.js
assets/noto-sans-devanagari-devanagari-wght-normal-DWI_L5Cu.woff2
assets/noto-sans-devanagari-latin-ext-wght-normal-CLKupaxV.woff2
assets/noto-sans-devanagari-latin-wght-normal-Bxfr4UJH.woff2
assets/noto-sans-kannada-kannada-wght-normal-DuWB-P8d.woff2
assets/noto-sans-kannada-latin-ext-wght-normal-MdBBF41x.woff2
assets/noto-sans-kannada-latin-wght-normal-DeLZCgFT.woff2
index.html
ui_version.json
=== ui_version.json ===
{

```

```
"spa_repo": "https://github.com/avidullu/khelsutra.git",
"spa_commit": "46e0244b0c862a4a9d845b8947843b5dfa1f2d9f",
"engine_api_version": 1,
"content_hash": "fad6253629831592b11dc59a1b4546af8a1fa96cc13c498a9255b47dcf0d04ff",
"built_at": "2026-06-21T17:35:18Z"
}
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

447. assistant (2026-06-21T17:35:27.187Z)

Bundler works ? `backend/ui/` has the SPA + provenance (`spa\_commit=46e0244`, `engine\_api\_version=1`). Now wire FastAPI to serve it single-origin. Let me read the current static-mount block + `/api/version` + the index import.

448. user (2026-06-21T17:35:29.176Z)

```
33:from fastapi.staticfiles import StaticFiles
34:from fastapi.responses import FileResponse, HTMLResponse, JSONResponse
128:_STATIC_DIR = os.path.join(os.path.dirname(os.path.abspath(__file__)), "api", "static")
130:    os.makedirs(_STATIC_DIR, exist_ok=True)
135:app.mount("/static", StaticFiles(directory=_STATIC_DIR), name="static")
137:@app.get("/", response_class=HTMLResponse)
138:def serve_index():
139:    index_path = os.path.join(_STATIC_DIR, "index.html")
236:API_VERSION = 1
239:def _engine_version() -> str:
248:def api_version():
250:    return {"engine_version": _engine_version(), "api_version": API_VERSION}
272:    "engine_version": _engine_version(),
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

449. user (2026-06-21T17:35:33.126Z)

```
126    # arbitrary directory finds the shipped assets instead of creating a stray
./backend/api/static.
127    # (The legacy in-engine dashboard; the real SPA gets bundled here in P1a.)
128    _STATIC_DIR = os.path.join(os.path.dirname(os.path.abspath(__file__)), "api", "static")
129    try:
130        os.makedirs(_STATIC_DIR, exist_ok=True)
131    except OSError:
132        # The dir ships with the package (a no-op create); guard against a read-only install
133        # tree (system-installed Python) since exist_ok suppresses FileExistsError, not
PermissionError.
134        pass
135    app.mount("/static", StaticFiles(directory=_STATIC_DIR), name="static")
136
137    @app.get("/", response_class=HTMLResponse)
138    def serve_index():
139        index_path = os.path.join(_STATIC_DIR, "index.html")
140        if os.path.exists(index_path):
141            with open(index_path, "r", encoding="utf-8") as f:
142                return f.read()
143        return "<h3>Sports Highlight Indexer Server is Running. Please create index.html in
backend/api/static/</h3>"
144
145    # Global variables initialized at startup
```

450. assistant (2026-06-21T17:36:15.871Z)

SPA is query-param based (no client-side path routing), so serving `index.html` at `/` + an `/assets` mount suffices ? no greedy catch-all needed (which would also be a route-ordering hazard). Wiring the bundled-UI serving + `json` import + the `/api/version` UI field.

451. user (2026-06-21T17:36:24.322Z)

The file C:\Users\avidu\Projects\bhi-wt-p1b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

452. user (2026-06-21T17:36:31.095Z)

The file C:\Users\avidu\Projects\bhi-wt-p1b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

453. assistant (2026-06-21T17:36:35.917Z)

Add the `import json` (used by `\_bundled\_ui\_version`).

454. user (2026-06-21T17:36:37.302Z)

The file C:\Users\avidu\Projects\bhi-wt-p1b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

455. assistant (2026-06-21T17:37:02.321Z)

Now the contract test (incl. the skew guard: the snapshot records the engine `API\_VERSION` it was built against).

456. assistant (2026-06-21T17:37:10.691Z)

```
"""The bundled SPA is served single-origin + the SPA<->engine contract (PACKAGING_PLAN.md P1b)."""
```

```
import json
import os
```

```
import pytest
```

```
pytest.importorskip("httpx")
from fastapi.testclient import TestClient
```

```
import backend.main as m
```

```
client = TestClient(m.app)
UI_DIR = os.path.join(os.path.dirname(os.path.abspath(m.__file__)), "ui")
```

```
def test_bundled_ui_snapshot_is_committed():
    assert os.path.isfile(os.path.join(UI_DIR, "index.html")), (
        "backend/ui/index.html missing ? the SPA snapshot must be committed (run scripts/bundle_ui.sh).")
    assert os.path.isfile(os.path.join(UI_DIR, "ui_version.json"))
```

```
def test_serves_spa_at_root():
    r = client.get("/")
    assert r.status_code == 200
    # The bundled index.html references its single-origin /assets bundle (Vite base='/').
    assert "/assets/" in r.text
```

```
def test_serves_spa_assets_byte_for_byte():
    assets = os.path.join(UI_DIR, "assets")
    js = next(f for f in os.listdir(assets) if f.endswith(".js"))
    r = client.get(f"/assets/{js}")
    assert r.status_code == 200
    with open(os.path.join(assets, js), "rb") as f:
        assert r.content == f.read() # served verbatim (mimetype varies by OS, so compare bytes)
```

```

def test_api_version_reports_bundled_ui():
    body = client.get("/api/version").json()
    assert body["api_version"] == m.API_VERSION
    assert body["ui"] is not None
    assert body["ui"]["spa_commit"] # provenance recorded

def test_spa_engine_contract_version_matches():
    """Skew guard: the committed snapshot records the engine API_VERSION it was built against.
    If
    someone bumps API_VERSION without re-running scripts/bundle_ui.sh, this FAILS ? forcing a
    re-bundle so the shipped SPA can't silently target a stale API contract."""
    with open(os.path.join(UI_DIR, "ui_version.json"), encoding="utf-8") as f:
        uiv = json.load(f)
    assert uiv["engine_api_version"] == m.API_VERSION, (
        f"bundled SPA built against engine api_version {uiv['engine_api_version']} but engine is
    now "
        f"{m.API_VERSION} ? re-run scripts/bundle_ui.sh and commit backend/ui/."
    )

```

457. user (2026-06-21T17:37:10.712Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-plb\tests\test_ui_bundle.py (file state is current in your context ? no need to Read it back)

458. assistant (2026-06-21T17:37:16.413Z)

Now verify `backend/ui/` lands in the wheel + run the new tests + lint.

459. user (2026-06-21T17:38:15.700Z)

```

=== build wheel, check backend/ui inclusion ===
* Building wheel...
Successfully built badminton_highlight_indexer-0.1.0-py3-none-any.whl
backend/ui files in wheel: 10
has index.html: True
has ui_version.json: True
sample: ['backend/ui/index.html', 'backend/ui/ui_version.json', 'backend/ui/assets/index-
FJh_kElx.css']
=== tests ===
..... [100%]
43 passed in 12.88s
=== ruff (script + main + test) ===
All checks passed!
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

460. assistant (2026-06-21T17:38:51.688Z)

The wheel auto-includes all 10 `backend/ui/` files (no force-include needed), 43 tests pass, lint clean. Since I'm committing built assets (minified JS/CSS + woff2 fonts), I must prevent autocrlf from mangling them (the known CRLF-drift footgun). Let me check `.gitattributes`.

461. user (2026-06-21T17:38:53.506Z)

```

=== .gitattributes ===

```

Golden Regression Fixtures (M1) ? pin LF line endings so the committed fixtures are byte-stable across OSes despite core.autocrlf. The characterization gate parse-compares (never byte-compares) JSON, but a stable on-disk form keeps `git diff` reviewable and fixtures reproducible. See docs/GOLDEN_REGRESSION_FIXTURES.md §4d/§4e.


```
eval_baselines/fixtures/** text eol=lf
```

Shell scripts must stay LF so they run on Linux boxes (the DigitalOcean bootstrap) regardless of a contributor's core.autocrlf. A CRLF shebang (`#!/usr/bin/env bash\r`) fails on Linux. See deploy/digitalocean/.

```
*.sh text eol=lf
=== would git mangle the committed JS/CSS? (autocrlf) ===
true
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

462. assistant (2026-06-21T17:39:09.306Z)

`core.autocrlf=true` ? so I must mark the vendored SPA build output as binary (`-text`) to keep it byte-exact cross-platform. Adding that.

463. user (2026-06-21T17:39:13.461Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

464. user (2026-06-21T17:39:17.188Z)

```
1      # Golden Regression Fixtures (M1) ? pin LF line endings so the committed fixtures are
2      # byte-stable across OSes despite core.autocrlf. The characterization gate parse-
compares
3      # (never byte-compares) JSON, but a stable on-disk form keeps `git diff` reviewable and
the
4      # fixtures reproducible. See docs/GOLDEN_REGRESSION_FIXTURES.md §4d/§4e.
5      eval_baselines/fixtures/** text eol=lf
6
7      # Shell scripts must stay LF so they run on Linux boxes (the DigitalOcean
8      # bootstrap) regardless of a contributor's core.autocrlf. A CRLF shebang
9      # (`#!/usr/bin/env bash\r`) fails on Linux. See deploy/digitalocean/.
10     *.sh text eol=lf
11
```

465. user (2026-06-21T17:39:22.713Z)

The file C:\Users\avidu\Projects\bhi-wt-p1b\gitattributes has been updated successfully. (file state is current in your context ? no need to Read it back)

466. assistant (2026-06-21T17:39:26.536Z)

Now the full suite + coverage.

467. user (2026-06-21T17:41:41.307Z)

```
backend\main.py                529    139    74%
TOTAL                          9836   1464    85%
Required test coverage of 70% reached. Total coverage: 85.12%
1295 passed, 7 skipped, 7 warnings in 129.79s (0:02:09)
  RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

468. assistant (2026-06-21T17:41:58.213Z)

Full suite green (1295 passed), coverage 85.12%. This is a meaty cross-repo PR (vendored assets + new serving + bundler) ? focused adversarial review before commit.

469. user (2026-06-21T17:45:45.301Z)

All uploads routes are `/api/upload/\*` prefixed ? no catch-all, no shadowing of `/` or

`/assets`. The `{upload\_id}` and `{index}` path params are scoped under `/api/upload/`, so they cannot intercept SPA asset requests.

I now have complete file:line evidence for all 7 review areas. Let me compile the findings.

Adversarial review ? PR P1b (bundle khelsutra SPA into engine wheel, single-origin)

[nit] backend/main.py:67 / scripts/bundle_ui.sh:67 ? `content\_hash` is path-dependent, not a reproducible content fingerprint ? `CONTENT\_HASH="\$(find "\$UI\_DEST" ... | xargs sha256sum | sha256sum ...)"` runs `find` over the *absolute* `\$UI\_DEST` path, and the outer `sha256sum` hashes a stream that embeds those absolute paths (verified: it reproduces `fad6253` only with the literal `/c/Users/avidu/Projects/bhi-wt-p1b/...` prefix; a relative-path recompute gives `ac2a3d6`). So re-bundling identical bytes from any other checkout path/machine/CI yields a different hash. It's recorded as provenance but isn't a verifiable fingerprint, and no test asserts it. ? Fix: hash with paths relativized to `\$UI\_DEST` (e.g. `cd "\$UI\_DEST" && find . -type f ! -name ui\_version.json | LC\_ALL=C sort | xargs sha256sum | sha256sum`), or drop the field if it's only decorative.

[nit] scripts/bundle_ui.sh:66 ? `API\_VERSION` extraction is brittle to a type annotation ? the anchor `^API\_VERSION\s\*=` matches today's `API\_VERSION = 1` (verified ? `1`), but a future `API\_VERSION: int = 1` would not match ? `ENGINE\_API\_VERSION` empty ? json written as `engine\_api\_version: null`. This fails *safe* (the skew test then asserts `null == 1` and fails loudly, prompting a re-bundle), so it's not a guard hole ? just a footgun. ? Fix: loosen to `^API\_VERSION\b[^\s]\*=` or note the assumption.

[minor] backend/ui/assets/*.js + *.woff2 ? minified third-party JS (React etc.) + OFL Noto fonts are committed into the MIT engine repo, and the license-scan lane does not cover them ? `.github/workflows/ci.yml:142` runs `pip-licenses` (Python-only); it cannot see the JS/font blobs. No `LICENSE`/`NOTICE` is vendored alongside the bundle (verified: none under `backend/ui/`). Given the project's hard MIT/Apache/BSD-only guardrail for *all* deps, and that OFL/MIT carry attribution obligations, this warrants an explicit note in the PR (and ideally a vendored third-party-notices file from the SPA build). Not a serving/correctness blocker. ? Fix: PR should NOTE the JS/font licensing posture; consider vendoring the SPA's generated licenses file into `backend/ui/`.

Verified clean (no issue):

- **Serving correctness** ? `serve\_index` (main.py:151) reads `\_UI\_DIR/index.html`; `/assets` mount (main.py:147) is guarded by `os.path.isdir(\_UI\_ASSETS\_DIR)` (146), registered at import only when the dir exists. No `/assets`-prefixed API route exists ? no conflict; `/static` (136) is a distinct prefix. Fallback when `backend/ui` absent: `\_UI\_DIR` ? `\_STATIC\_DIR` (144), legacy `index.html` references `/static/\*` served by the always-present `/static` mount (verified functional, no crash); `\_UI\_ASSETS\_DIR` becomes a non-existent `static/assets` so `/assets` is simply not mounted. Traversal: Starlette 1.0.0 `StaticFiles` normalizes and 404s `../` escapes.
- **Route order / shadowing** ? no `{path:path}` catch-all anywhere (main.py + `backend/api/uploads.py` router routes are all `/api/upload/\*`). Serving `index.html` at `/` (150) does not intercept `/api/\*` or `/static`.
- **Bundler robustness** ? `set -euo pipefail` (sh:17); option-3 uses `curl -fsSL` + `tar -xzf` into a `mktemp -d` with `trap ? EXIT` cleanup; `stage\_dist` validates `index.html` presence and does `rm -rf` + `mkdir -p` + `cp -r "\$dist"/.` into a clean dest (correct). Untrusted-tarball traversal is the standard `tar` caveat but the source is the owner's own release artifact (acceptable). CI never bundles ? no `.github/` reference to `bundle\_ui`/`npm`/`backend/ui` (verified); the snapshot is committed.
- **Skew guard** ? genuine: `test\_spa\_engine\_contract\_version\_matches` (test_ui_bundle.py:46) loads the committed `ui\_version.json` and asserts `engine\_api\_version == m.API\_VERSION`; bumping `API\_VERSION` without re-bundling makes them differ ? FAIL. No spurious-pass path.
- **CRLF/binary** ? `.gitattributes:15` backend/ui/** -text resolves to `text: unset` for `.html/.js/.json/.woff2` (verified via `git check-attr`); repo has `core.autocrlf=true`, so this rule is *necessary* and correct ? committed bytes will match Vite output. `ui\_version.json` is also covered by *-text* (it's generated, never hand-edited, so no LF normalization needed ? fine).
- **/api/version robustness** ? `\_bundled\_ui\_version` (main.py:159) reads `\_BUNDLED\_UI\_DIR/ui\_version.json` and returns `None` on `(OSError, ValueError)` (covers missing

```

file + parse error); `json` imported (main.py:4), `Optional/Dict/Any` imported (main.py:32).
`/api/version` still 200s without a bundle.
- Wheel packaging ? built the wheel from the current (untracked `backend/ui/`) tree: all 10
UI files including the woff2 fonts are present (`backend/ui/...`). No force-include needed;
claim holds.
- Targeted tests ? `tests/test_ui_bundle.py`: 5 passed.

One process flag (not a code issue): the entire change is uncommitted working-tree state on
`feat/packaging-p1b-bundle-ui` (HEAD is still P1a #300); `backend/ui/` is untracked,
`.gitattributes`/`backend/main.py` modified-not-staged. All verification above reflects working-
tree bytes ? re-confirm after `git add` that the `-text` attribute lands the bytes unmangled (it
will, since `.gitattributes` is in-tree). Did not re-run the full 1295-test suite / ruff / mypy
/ coverage (prompt states these were already verified green at 85.12%).

Verdict: No blockers or majors ? serving, fallback, route-ordering, and skew-guard logic are
correct and well-guarded; address the JS/font licensing note (minor) and optionally the two
bundler nits before merge.
agentId: a9ce9fd6661c86dc0 (use SendMessage with to: 'a9ce9fd6661c86dc0' to continue this agent)
<usage>subagent_tokens: 66941
tool_uses: 26
duration_ms: 208537</usage>

```

470. assistant (2026-06-21T17:46:27.443Z)

Review: no blockers/majors. Three worth folding: vendor a **third-party-notices** file for the committed JS/fonts (the python license-scan can't see them), make the **content_hash** reproducible (it embedded absolute paths), and **loosen the `API\_VERSION` grep**. Updating the bundler.

471. user (2026-06-21T17:46:36.539Z)

The file C:\Users\avidu\Projects\bhi-wt-p1b\scripts\bundle_ui.sh has been updated successfully. (file state is current in your context ? no need to Read it back)

472. user (2026-06-21T17:46:50.542Z)

```

assets/noto-sans-kannada-latin-wght-normal-DeLZCgFT.woff2
index.html
ui_version.json
=== ui_version.json (reproducible hash) ===
{
  "spa_repo": "https://github.com/avidullu/khelsutra.git",
  "spa_commit": "46e0244b0c862a4a9d845b8947843b5dfa1f2d9f",
  "engine_api_version": 1,
  "content_hash": "ac2a3d672c016d57c3e237f17b9251b22a0c0a084d488906b586aeecf4a666a3",
  "built_at": "2026-06-21T17:46:43Z"
}
=== NOTICE present ===

```

Third-party notices ? vendored KhelSutra SPA (`backend/ui/`)

```

This directory is the BUILT KhelSutra web SPA (MIT), vendored from
https://github.com/avidullu/khelsutra.git @ 46e0244b0c862a4a9d845b8947843b5dfa1f2d9f by
=== re-run tests + ruff ===
..... [100%]
5 passed in 5.37s
Found 207 errors.
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

473. assistant (2026-06-21T17:47:09.676Z)

Hash is now reproducible (`ac2a3d6?`), NOTICE vendored, tests pass. (The "207 errors" was just ruff wrongly parsing the `.sh` as Python ? CI's ruff only lints `.py`.) Let me syntax-check the script, stage explicit paths, and confirm the **-text** attribute lands on the assets.

474. user (2026-06-21T17:47:15.630Z)

```
=== bash syntax ===
OK
warning: in the working copy of 'tests/test_ui_bundle.py', LF will be replaced by CRLF the next
time Git touches it
=== -text attr lands on committed assets? ===
backend/ui/index.html: text: unset
backend/ui/assets/index-ZjpoX6Fg.js: text: unset
backend/ui/assets/noto-sans-kannada-kannada-wght-normal-DuwB-P8d.woff2: text: unset
=== staged summary ===
M .gitattributes
M backend/main.py
A backend/ui/THIRD_PARTY_NOTICES.md
A backend/ui/assets/index-FJh_kElx.css
A backend/ui/assets/index-ZjpoX6Fg.js
A backend/ui/assets/noto-sans-devanagari-devanagari-wght-normal-DWI_L5Cu.woff2
A backend/ui/assets/noto-sans-devanagari-latin-ext-wght-normal-CLKupaxV.woff2
A backend/ui/assets/noto-sans-devanagari-latin-wght-normal-Bxfr4UJH.woff2
A backend/ui/assets/noto-sans-kannada-kannada-wght-normal-DuwB-P8d.woff2
A backend/ui/assets/noto-sans-kannada-latin-ext-wght-normal-MdBBF41x.woff2
A backend/ui/assets/noto-sans-kannada-latin-wght-normal-DeLZCgFT.woff2
A backend/ui/index.html
A backend/ui/ui_version.json
A scripts/bundle_ui.sh
A tests/test_ui_bundle.py
=== count staged backend/ui files ===
11
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

475. assistant (2026-06-21T17:47:22.503Z)

```
`-text` lands on the assets (`text: unset` ? byte-exact), bash syntax OK, 11 `backend/ui/` files
+ script + main + gitattributes + test staged. Committing P1b.
```

476. user (2026-06-21T17:47:38.046Z)

```
=== HEAD ===
411e19f feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel (default-
OFF)
branch 'feat/packaging-p1b-bundle-ui' set up to track 'origin/feat/packaging-p1b-bundle-ui'.
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch]      feat/packaging-p1b-bundle-ui -> feat/packaging-p1b-bundle-ui
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

477. user (2026-06-21T17:47:58.694Z)

```
https://github.com/Khelsutra/badminton-highlight-indexer/pull/302
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

478. assistant (2026-06-21T17:48:02.834Z)

```
**PR #302 opened: [feat(pkg): P1b ? bundle the SPA single-
origin](https://github.com/Khelsutra/badminton-highlight-indexer/pull/302).** Watching CI (the
built-wheel lane will now exercise the bundled UI inclusion).
```

479. user (2026-06-21T17:50:29.361Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    21s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27912518268/job/82591995305
Built wheel smoke (install + launch)             pass    1m6s  https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27912518268/job/82591995285
Dependency license scan (no copyleft in runtime deps) pass    34s
```

```
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995300
Full test suite (offline, mocked) pass 2m20s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995331
Golden fixtures cross-OS (macos-latest) pass 21s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995312
Golden fixtures cross-OS (ubuntu-latest) pass 26s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995376
Golden fixtures cross-OS (windows-latest) pass 1m29s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995297
Mypy (lenient baseline) pass 36s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995294
Ruff lint pass 8s https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995287
Syntax sweep + import smoke (no GPU stack) pass 21s
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27912518268/job/82591995305
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

480. assistant (2026-06-21T17:50:32.864Z)

All 9 lanes green. Merging P1b and cleaning up.

481. user (2026-06-21T17:50:42.887Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
be0c434..4a06b8b master -> origin/master
=== master ===
4a06b8b feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel (default-OFF) (#302)
=== PR #302 ===
MERGED
=== open PRs (mine) ===
302 feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel
feat/packaging-p1b-bundle-ui MERGED 2026-06-21T17:47:58Z
```

482. assistant (2026-06-21T17:50:56.897Z)

P1b ****merged**** (master `4a06b8b`) ? ****P1** is now complete: the engine serves the real KhelSutra SPA single-origin from the wheel.****** Updating the tracker + memory.

483. user (2026-06-21T17:51:04.349Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

484. user (2026-06-21T17:51:17.625Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

485. assistant (2026-06-21T17:51:49.502Z)

Clean checkpoint ? ****P0** + **P1** both complete, 7 PRs merged this session, 0 open PRs, no VMS (\$0 spend), all worktrees cleaned.****** Here's the full picture.

What shipped this session

****The plan**** ? docs/PACKAGING_PLAN.md: one profile-selected build, cross-platform webserver (no native component), LIGHT-first with a WASB-complete north star (your 01/02/bundling calls baked in).

****P0 ? engine prerequisites (5 PRs):****

```
| PR | |
|---|---|
| [#293](https://github.com/Khelsutra/badminton-highlight-indexer/pull/293) | App-data HOME
(`RALLY_APP_HOME`) + fixed the doctor?server config footgun |
| [#295](https://github.com/Khelsutra/badminton-highlight-indexer/pull/295) | Redacted the
unauthenticated `/api/config` |
| [#296](https://github.com/Khelsutra/badminton-highlight-indexer/pull/296) | Config-driven bind
+ DNS-rebinding Host guard |
| [#297](https://github.com/Khelsutra/badminton-highlight-indexer/pull/297) | Declared `scipy` +
license-scan & built-wheel CI lanes |
| [#298](https://github.com/Khelsutra/badminton-highlight-indexer/pull/298) | Single-instance
lock (no 2nd process corrupts job state) |
```

****P1 ? the UI is now in the box (2 PRs):****

```
| PR | |
|---|---|
| [#300](https://github.com/Khelsutra/badminton-highlight-indexer/pull/300) | `/api/version` +
`/api/capabilities` + `/api/videos` + `/api/reels` (console/wizard endpoints) |
| [#302](https://github.com/Khelsutra/badminton-highlight-indexer/pull/302) | **Bundled the
KhelSutra SPA single-origin into the wheel** ? `scripts/bundle_ui.sh` + committed `backend/ui/`
snapshot + skew-guard contract test |
```

****Net result:**** `pip install` the engine wheel ? it ships the real KhelSutra UI and serves it at `localhost` single-origin, finds all its state in a resolvable app-home, refuses unsafe exposure, and has CI guarding the wheel + licenses. That's the engine fully ****bundle-ready****.

Every PR: isolated worktree off `origin/master` (which moved ~9× under me ? no collisions), explicit-path commits, adversarial review folded, full suite + 9-lane CI green, coverage held ~85%. The cross-repo SPA bundling uses your ****option 1**** (committed snapshot + `bundle\_ui.sh`), built to also consume a ****release artifact (option 3)**** later.

Why I'm checkpointing here

This is a clean, meaningful boundary (the engine half of "batteries-included" is done). ****P2**** is the next tranche ? a cross-platform launcher (`indexer --app` ? boot + open browser) ****plus**** a first-run wizard. The launcher is engine-only and unblocked; the ****wizard UI is SPA work in the separate `khelsutra` repo**** again (like P1b), so it's worth your steer before I start ? and I've been running autonomously through a lot, so a checkpoint is prudent.

Want me to:

1. ****Continue into P2's launcher**** (`indexer --app`, engine-only, unblocked) now, and bring the wizard (cross-repo) to you separately, or
2. ****Pause here**** for you to review the merged P0/P1 work first, or
3. Something else (e.g., jump to ****P3 GPU/cloud validation**** using the VM budget)?

486. user (2026-06-21T17:56:23.605Z)

There have been more activity on the repo now. Can you sync to remote head and ground you recommendation of next steps? Ideally I would want you to continue the P2 and onward but want to ground at a checkpoint first.

487. assistant (2026-06-21T17:56:37.413Z)

Good call ? let me sync and re-ground before recommending P2. Fetching and checking what's landed since my last merge (`4a06b8b`).

488. user (2026-06-21T17:56:42.900Z)

=== current main-checkout branch (sibling-owned) ===
serving/exposure-safety-preflight

```

6fdd1d4 feat(serving): exposure-safety boot posture ? fail fast on a half-configured exposed
engine (default-OFF)

=== origin/master HEAD ===
4a06b8b feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel (default-
OFF) (#302)

=== NEW on origin/master since my P1b merge (4a06b8b) ===

=== my packaging PRs still present on master? ===
4a06b8b feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel (default-
OFF) (#302)
be0c434 feat(pkg): P1a ? net-new console/wizard endpoints: /api/version, /api/capabilities,
/api/videos, /api/reels (#300)
c039b90 feat(pkg): P0e ? single-instance lock so a 2nd process can't corrupt in-flight job state
(default-OFF) (#298)
e8bad40 feat(pkg): P0d ? declare scipy (undeclared serving dep) + license-scan & built-wheel CI
lanes (#297)
3895963 feat(pkg): P0c ? config/profile-driven bind address + DNS-rebinding Host-header guard
(default-OFF) (#296)
fc78fce feat(pkg): P0b ? redact secret-shaped fields from the unauthenticated GET /api/config
(#295)
82a28f4 feat(pkg): P0a ? resolved app-data HOME (RALLY_APP_HOME) for config/.env/DB/output
(default-OFF) (#293)

```

489. assistant (2026-06-21T17:56:59.282Z)

`origin/master` is still at my P1b (`4a06b8b`) ? nothing new merged. But the main checkout is on an **unmerged sibling branch** `serving/exposure-safety-preflight` (exposure-safety boot posture ? which overlaps my P0c startup checks). So the "activity" is likely in-flight PRs and/or the SPA repo. Let me ground on all three.

490. user (2026-06-21T17:57:05.715Z)

```

=== OPEN PRs on the engine repo (all authors) ===

=== khelsutra (SPA) head ? did it advance past my pinned 46e0244? ===
f7e2bd5 chore(site): add robots.txt + a "crawling / bot protection" runbook for khelsutra.guru
(#71)
my bundle pinned: 46e0244

=== sibling branch 'serving/exposure-safety-preflight' vs master: what does it touch? ===
backend/api/auth.py | 6 +-
backend/config/startup.py | 113 ++++++-----
backend/main.py | 14 +--
docs/COMPUTE_DECOUPLED_SERVING/README.md | 3 +-
tests/test_edge_auth.py | 9 +++
tests/test_startup_checks.py | 132 ++++++-----
6 files changed, 254 insertions(+), 23 deletions(-)

```

491. assistant (2026-06-21T17:57:29.659Z)

No open engine PRs. Two relevant findings: (1) **khelsutra advanced** `46e0244` ? f7e2bd5` (I bundled `46e0244`), and (2) a sibling branch `serving/exposure-safety-preflight` is reworking `startup.py`/`main.py` (building **on** my merged P0c). Let me check whether the SPA's `web/` actually changed (would my bundle be stale?) and peek at the sibling's startup-check direction.

492. user (2026-06-21T17:57:33.767Z)

```

=== khelsutra commits 46e0244..f7e2bd5 ===
f7e2bd5 chore(site): add robots.txt + a "crawling / bot protection" runbook for khelsutra.guru
(#71)
bea8a2a docs(appliance): flip T11 retention sweep to ? ? tool shipped (engine #299), scheduling
tracked (#301) (#69)

```

=== did the SPA web/ change since my bundle? (web/ files touched) ===

=== sibling exposure branch: new EXPOSED_* check codes (overlap w/ my P0c?) ===

```
EXPOSED_AUTH_EXTRA_MISSING
EXPOSED_EDGE_AUTH_INCOMPLETE
EXPOSED_NO_VIDEO_ROOT
```

493. assistant (2026-06-21T17:58:51.210Z)

Grounding complete. Here's the picture:

```
**State is stable and my work is intact:**
- `origin/master` = my P1b (`4a06b8b`); **all 7 P0/P1 PRs present**, **0 open PRs**.
- **The bundled SPA is current** ? khelsutra only moved on `site/+docs` (`#69`, `#71`); **zero
`web/` changes** since my pinned `46e0244`, so `backend/ui/` doesn't need a re-bundle.
- The engine surface P2 builds on (bundled UI serving, `/api/capabilities`, the `main()`
entrypoint) is exactly as I left it.
```

```
**One coordination note (not blocking):** a sibling WIP branch `serving/exposure-safety-
preflight` (no PR yet) is reworking `startup.py`/`main.py` exposure checks ? and its
`startup.py` appears to **lack my P0c `EXPOSED_HOST_ALLOWLIST_LOOPBACK`** (likely branched
pre-P0c). When it opens a PR it'll need rebasing onto current master so that check isn't
dropped. I'll branch P2 off `origin/master` + stage explicit paths so I don't collide, and flag
this for whoever reviews that PR.
```

```
**Recommendation (grounded):** P2 next = the **launcher (`indexer --app`)** as the headline
batteries-included PR ? engine-only, unblocked, and the thing that makes it *feel* like an app
(boot server ? open browser to the bundled SPA). Then the ffmpeg resolver (P2a) and DB-upgrade
safety (P2d), with the cross-repo **first-run wizard** (P2c, SPA work in khelsutra) coordinated
like P1b.
```

Continuing now with the launcher. Setting up the worktree and reading the current `main()`
server-start.

494. user (2026-06-21T17:58:55.093Z)

HEAD is now at 4a06b8b feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine
wheel (default-OFF) (#302)

=== head ===

4a06b8b feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel (default-
OFF) (#302)

=== --server/--port args + the server-start block ===

```
800:     parser.add_argument("--server", action="store_true", help="Run the FastAPI local API
server")
801:     parser.add_argument("--port", type=int, default=8000, help="Port to run the FastAPI
server on")
913:     if args.video_path and not args.server:
994:     if args.server or not args.video_path:
995:         bind_host = _resolve_bind_host(global_config, args.host) # P0c: config/profile-
driven (default 127.0.0.1)
996:         print(f"[SERVER] Starting local FastAPI server on
http://{bind_host}:{args.port}...")
1006:         uvicorn.run(app, host=bind_host, port=args.port)
```

495. user (2026-06-21T17:58:58.747Z)

```
994         if args.server or not args.video_path:
995             bind_host = _resolve_bind_host(global_config, args.host) # P0c: config/profile-
driven (default 127.0.0.1)
996             print(f"[SERVER] Starting local FastAPI server on
http://{bind_host}:{args.port}...")
997             if bind_host not in ("127.0.0.1", "localhost", ":::1"):
998                 print(f"[SERVER] ? Binding {bind_host} exposes the API on the network ?
ensure an auth ")
```



```

999             f"gate (edge_auth_enabled / a reverse proxy) is in front, and set "
1000             f"RALLY_ALLOWED_HOSTS to your public host(s) (the Host-header guard
defaults to "
1001             f"loopback-only and will 400 forwarded requests otherwise).")
1002         # Resume self-scheduled work orphaned by the restart (EXECUTION_POLICY.md P3):
the
1003         # in-process wake-timer for any parked 'waiting' job died with the old process,
so
1004         # kick one drain ? it reclaims due jobs and re-arms the wake-timer for the rest.
1005         _get_executor().submit(_drain_due_jobs)
1006         uvicorn.run(app, host=bind_host, port=args.port)
1007
1008     if __name__ == "__main__":
1009         # `python -m backend.main` runs THIS file as `__main__`, but `job_worker`/`uploads`
resolve the

```

496. assistant (2026-06-21T17:59:31.837Z)

Clear. Implementing `indexer --app`: a pure-Python launcher (force localhost, free-port fallback, open browser when healthy, `RALLY\_NO\_BROWSER` escape for tests/headless). Adding imports + helpers.

497. user (2026-06-21T17:59:33.404Z)

The file C:\Users\avidu\Projects\bhi-wt-p2b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

498. user (2026-06-21T17:59:41.922Z)

```

113: def _resolve_bind_host(config: Optional[AppConfig] = None, cli_host: Optional[str] = None)
-> str:
126:     return "127.0.0.1"
797: def main():
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

499. user (2026-06-21T17:59:46.080Z)

```

124         if getattr(sec, "edge_auth_enabled", False):
125             return "0.0.0.0" # behind a Cloudflare-Access-style edge gate (M9) ? bind all
interfaces
126             return "127.0.0.1"
127
128     # Serves static frontend files directly (now under api/ per approved structure).

```

500. user (2026-06-21T17:59:58.234Z)

The file C:\Users\avidu\Projects\bhi-wt-p2b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

501. assistant (2026-06-21T18:00:02.467Z)

Now add the `--app` arg, guard the CLI branch, and split the server-start block.

502. user (2026-06-21T18:00:05.130Z)

The file C:\Users\avidu\Projects\bhi-wt-p2b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

503. user (2026-06-21T18:00:09.364Z)

```

967         logger.warning(f"[JOBS] Swept {swept} stale job(s) from a previous run ?
failed.")
968

```

```
969         # CLI processing mode (no web UI needed)
970         if args.video_path and not args.server:
971             print(f"[CLI] Processing video file: {args.video_path}")
```

504. user (2026-06-21T18:00:15.170Z)

The file C:\Users\avidu\Projects\bhi-wt-p2b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

505. assistant (2026-06-21T18:00:18.966Z)

Now the server-start block ? split `--app` (friendly launcher, localhost-only) from plain `--server`.

506. user (2026-06-21T18:00:30.907Z)

The file C:\Users\avidu\Projects\bhi-wt-p2b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

507. assistant (2026-06-21T18:01:06.236Z)

Now tests ? unit (port helpers + browser-open logic) and a subprocess boot of `indexer --app` (with `RALLY\_NO\_BROWSER=1`) proving it serves health + the bundled SPA.

508. assistant (2026-06-21T18:01:21.957Z)

```
"""The friendly `indexer --app` launcher (PACKAGING_PLAN.md P2)."""
```

```
import os
import socket
import subprocess
import sys
import time
import urllib.error
import urllib.request
```

```
import pytest
```

```
import backend.main as m
```

```
REPO_ROOT = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
```

```
def _free_port() -> int:
    s = socket.socket()
    s.bind(("127.0.0.1", 0))
    p = int(s.getsockname()[1])
    s.close()
    return p
```

---- port helpers ----

```
def test_find_and_check_free_port():
    p = m._find_free_port("127.0.0.1")
    assert isinstance(p, int) and p > 0
    assert m._port_is_free("127.0.0.1", p) is True
```

```
def test_resolve_app_port_keeps_free_port():
    p = _free_port()
    assert m._resolve_app_port("127.0.0.1", p) == p
```

```

def test_resolve_app_port_falls_back_when_busy():
    s = socket.socket()
    s.bind(("127.0.0.1", 0))
    s.listen()
    busy = int(s.getsockname()[1])
    try:
        assert m._port_is_free("127.0.0.1", busy) is False
        alt = m._resolve_app_port("127.0.0.1", busy)
        assert alt != busy
        assert m._port_is_free("127.0.0.1", alt) is True
    finally:
        s.close()

```

---- open-browser-when-ready ----

```

class _FakeResp:
    status = 200
    def __enter__(self): return self
    def __exit__(self, *a): return False

def test_open_browser_opens_once_healthy(monkeypatch):
    monkeypatch.delenv("RALLY_NO_BROWSER", raising=False)
    monkeypatch.setattr(urllib.request, "urlopen", lambda *a, **k: _FakeResp())
    opened = {}
    monkeypatch.setattr(m.webbrowser, "open", lambda u: opened.setdefault("url", u) or True)
    assert m._open_browser_when_ready("http://127.0.0.1:1/", "http://127.0.0.1:1/api/health",
    timeout=2) is True
    assert opened["url"] == "http://127.0.0.1:1/"

def test_open_browser_suppressed_by_env(monkeypatch):
    monkeypatch.setenv("RALLY_NO_BROWSER", "1")
    monkeypatch.setattr(m.webbrowser, "open", lambda u: pytest.fail("must not open when
    suppressed"))
    assert m._open_browser_when_ready("http://x/", "http://x/api/health") is False

def test_open_browser_gives_up_when_never_ready(monkeypatch):
    monkeypatch.delenv("RALLY_NO_BROWSER", raising=False)

    def _boom(*a, **k):
        raise OSError("connection refused")

    monkeypatch.setattr(urllib.request, "urlopen", _boom)
    monkeypatch.setattr(m.webbrowser, "open", lambda u: pytest.fail("must not open a dead
    page"))
    assert m._open_browser_when_ready("http://x/", "http://x/api/health", timeout=0.4) is False

```

---- end-to-end: `indexer --app` boots the server + serves the bundled UI ----

```

def test_app_mode_boots_and_serves_bundled_ui(tmp_path):
    """`indexer --app` (browser suppressed) must start the server localhost-only and serve both
    /api/health and the bundled SPA at /."""
    port = _free_port()
    env = {
        **os.environ,
        "RALLY_NO_BROWSER": "1", # don't pop a browser in CI
        "PYTHONPATH": REPO_ROOT + os.pathsep + os.environ.get("PYTHONPATH", ""),
    }
    log = open(tmp_path / "app.log", "wb")

```

```

proc = subprocess.Popen(
    [sys.executable, "-m", "backend.main", "--app", "--port", str(port),
      "--output-dir", str(tmp_path / "out")],
    cwd=REPO_ROOT, env=env, stdout=log, stderr=subprocess.STDOUT,
)
base = f"http://127.0.0.1:{port}"
try:
    deadline = time.time() + 90
    healthy = False
    while time.time() < deadline:
        if proc.poll() is not None:
            pytest.fail(f"--app server exited early (code {proc.returncode}); see
{tmp_path/'app.log'}")
            try:
                with urllib.request.urlopen(f"{base}/api/health", timeout=2) as r:
                    if r.status == 200:
                        healthy = True
                        break
            except (urllib.error.URLError, ConnectionError, OSError):
                time.sleep(0.5)
    assert healthy, "the --app server never became healthy"
    with urllib.request.urlopen(f"{base}/", timeout=5) as r:
        body = r.read().decode("utf-8", "replace")
    assert "/assets/" in body # the bundled SPA index.html (single-origin asset refs)
finally:
    proc.terminate()
    try:
        proc.wait(timeout=10)
    except subprocess.TimeoutExpired:
        proc.kill()
    log.close()

```

509. user (2026-06-21T18:01:21.971Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p2b\tests\test_app_launcher.py
(file state is current in your context ? no need to Read it back)

510. user (2026-06-21T18:01:40.330Z)

```

..... [100%]
7 passed in 11.25s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

511. assistant (2026-06-21T18:02:02.909Z)

All 7 launcher tests pass ? including the subprocess boot of `indexer --app` serving the bundled SPA. Now lint/type + full suite.

512. user (2026-06-21T18:04:55.395Z)

```

=== RUFF ===
All checks passed!
=== MYPY ===
backend\providers\__init__.py:7: note: By default the bodies of untyped functions are not
checked, consider using --check-untyped-defs [annotation-unchecked]
Success: no issues found in 1 source file
=== FULL SUITE + COV ===
backend\main.py          574    150    74%
TOTAL                   9881   1475    85%
Required test coverage of 70% reached. Total coverage: 85.07%
1302 passed, 7 skipped, 7 warnings in 132.33s (0:02:12)
RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

513. assistant (2026-06-21T18:05:04.471Z)

Full suite green (1302), lint/type clean. Quick focused review of the launcher (mode interactions + the daemon-thread browser open), then commit.

514. user (2026-06-21T18:06:41.665Z)

All confirmed. The health endpoint is auth-free and the poll uses `Host: 127.0.0.1:<port>` which passes the TrustedHostMiddleware default allowlist. Note one interaction worth flagging: if a user sets `RALLY\_ALLOWED\_HOSTS` without including `127.0.0.1`, the health poll would 400 ? but the default keeps `127.0.0.1`, and `--app` forces 127.0.0.1 bind, so the default path is safe. Setting `RALLY\_ALLOWED\_HOSTS` to an exclusive non-loopback set while using `--app` is a self-inflicted misconfiguration; the health poll would then never see 200 and the browser simply wouldn't open (no dead page). Acceptable.

Here are my findings.

****Adversarial review ? PR P2b (`indexer --app` launcher)****

****1. Mode interactions**** ? Clean. The CLI single-shot guard at `backend/main.py:970` (`if args.video\_path and not args.server and not args.app`) correctly excludes `--app`, so `--video-path --app` never runs single-shot processing and falls through to the server block at `:1051` (`if args.app or args.server or not args.video\_path`). `--app --server` ? `--app` branch wins (`:1052` checks `args.app` first) ? localhost-only, which is the safe precedence. No-args ? server (`not args.video\_path`). No combination runs two modes. The `--setup`/`--trim`/`--debug-clip` branches (`:872`, `:886`, `:918`) return early and require their own flags, independent of `--app`.

****2. Localhost invariant**** ? Confirmed. `:1055` hardcodes `bind\_host = "127.0.0.1"` for `--app`, fully bypassing `\_resolve\_bind\_host` / `--host` / `RALLY\_BIND\_HOST` / `edge\_auth\_enabled`. `--server` still honors the P0c profile via `\_resolve\_bind\_host(global\_config, args.host)` at `:1065`. Intended and safe ? the friendly app can never expose the network.

****3. Browser-open thread**** ? Solid. Daemon thread (`:1062-1063` `daemon=True`) ? won't block process exit or Ctrl+C. `RALLY\_NO\_BROWSER` honored at `:158`. Dead-page guard: the `while/else` at `:163-171` returns False without opening if health never returns 200 (no open on a server that fails to bind). `webbrowser.open` is wrapped in try/except (`:172-175`), exceptions swallowed. No race: it polls `/api/health` and only opens after a 200. Health URL is `url + "api/health" = "http://127.0.0.1:<port>/api/health" (url ends in `/`, `:1059`) ? `urllib` sends `Host: 127.0.0.1:<port>`, which passes TrustedHostMiddleware (port stripped ? `127.0.0.1`, in the default allowlist at `:100`). `/api/health` is auth-free (`:298`).

****4. Free-port (TOCTOU)**** ? Acceptable as documented. `\_resolve\_app\_port` (`:149-151`) checks `\_port\_is\_free` then uvicorn binds later. `\_port\_is\_free` (`:131`) binds without `SO\_REUSEADDR`; uvicorn binds with `SO\_REUSEADDR=1` ? a strictly-free port passes both, so the check is not falsely positive. `\_find\_free\_port` (`:142`) returns an OS-assigned ephemeral port (free at that instant). Window is best-effort; uvicorn surfaces a real bind failure loudly (docstring `:132-133`). The resolved `port` is used consistently for BOTH the browser URL (`:1059`) and `uvicorn.run` (`:1077`).

****5. Single-instance (P0e)**** ? Correct. `--app` reaches `acquire\_lock(db\_path + ".lock")` at `:946` and `fail\_stale\_jobs()` at `:965` (no early return between `:924` and the lock ? verified). A 2nd `--app` on the same `--output-dir` ? lock fails ? `[FATAL]` + `sys.exit(1)` at `:956-959`. Lock-infra failure (read-only dir) warns and proceeds without the guard (`:951`), not a crash.

****6. Resource/socket leaks**** ? None. Both `\_port\_is\_free` (`:134`) and `\_find\_free\_port` (`:144`) use `with socket.socket(...)` ? sockets always closed. `urllib.request.urlopen` uses `with` (`:165`). `webbrowser.open` exceptions swallowed (`:174`).

****[nit] backend/main.py:95-100**** ? If a user sets `RALLY\_ALLOWED\_HOSTS` to an exclusively non-loopback set *and* uses `--app`, the health poll's `Host: 127.0.0.1` would 400 and the browser would silently never open (no dead page, no crash). Self-inflicted misconfiguration and the

default path is safe; worth a one-line note that `--app` always needs loopback in the allowlist, but not blocking.

```
**Verdict:** Clean ? no blocker/major/minor issues; one nit. Ship it.  
agentId: a919b4e427b4a7016 (use SendMessage with to: 'a919b4e427b4a7016' to continue this agent)  
<usage>subagent_tokens: 55298  
tool_uses: 14  
duration_ms: 82351</usage>
```

515. assistant (2026-06-21T18:06:49.390Z)

Clean review ? one nit (an exotic self-inflicted `RALLY\_ALLOWED\_HOSTS` misconfig). Adding a one-line note, then committing.

516. user (2026-06-21T18:06:53.567Z)

The file C:\Users\avidu\Projects\bhi-wt-p2b\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

517. user (2026-06-21T18:07:12.380Z)

```
All checks passed!  
=== origin/master (still my P1b?) ===  
4a06b8b feat(pkg): P1b ? bundle the KhelSutra SPA single-origin into the engine wheel (default-OFF) (#302)  
=== status ===  
  M backend/main.py  
  ?? tests/test_app_launcher.py  
warning: in the working copy of 'tests/test_app_launcher.py', LF will be replaced by CRLF the next time Git touches it  
=== HEAD ===  
5dab7d5 feat(pkg): P2b ? friendly `indexer --app` launcher (boot localhost server + open the bundled UI)  
branch 'feat/packaging-p2b-launcher' set up to track 'origin/feat/packaging-p2b-launcher'.  
To https://github.com/Khelsutra/badminton-highlight-indexer.git  
  * [new branch]      feat/packaging-p2b-launcher -> feat/packaging-p2b-launcher  
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

518. user (2026-06-21T18:07:30.242Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/303>
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

519. assistant (2026-06-21T18:07:34.011Z)

```
**PR #303 opened: [feat(pkg): P2b ? `indexer --app`  
launcher](https://github.com/Khelsutra/badminton-highlight-indexer/pull/303).** Watching CI.
```

520. user (2026-06-21T18:10:23.253Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    24s  
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27913017890/job/82593316126  
Built wheel smoke (install + launch)            pass    1m7s   https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27913017890/job/82593316155  
Dependency license scan (no copleft in runtime deps) pass    35s  
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27913017890/job/82593316127  
Full test suite (offline, mocked)               pass    2m38s   https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27913017890/job/82593316138  
Golden fixtures cross-OS (macos-latest)         pass    16s     https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27913017890/job/82593316113  
Golden fixtures cross-OS (ubuntu-latest)        pass    21s  
https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27913017890/job/82593316113
```

```
indexer/actions/runs/27913017890/job/82593316117
Golden fixtures cross-OS (windows-latest)      pass    1m26s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913017890/job/82593316118
Mypy (lenient baseline) pass    29s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913017890/job/82593316147
Ruff lint      pass    9s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913017890/job/82593316108
Syntax sweep + import smoke (no GPU stack)      pass    24s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913017890/job/82593316126
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

521. assistant (2026-06-21T18:10:27.253Z)

All 9 lanes green. Merging P2b and cleaning up.

522. user (2026-06-21T18:10:35.841Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
4a06b8b..ca34787 master -> origin/master
=== master ===
ca34787 feat(pkg): P2b ? friendly `indexer --app` launcher (boot localhost server + open the
bundled UI) (#303)
=== PR #303 ===
MERGED
```

523. assistant (2026-06-21T18:10:57.857Z)

P2b ****merged**** (master `ca34787`) ? `indexer --app` is the headline batteries-included moment.
Let me update the tracker precisely (I shipped the ***launcher*** half of P2b; the uv/install-path
half remains) + memory.

524. user (2026-06-21T18:10:58.706Z)

```
98
99      **? P1 COMPLETE (P1a #300 + P1b #302) ? the engine now serves the real KhelSutra SPA
single-origin from the wheel, with the console/wizard endpoints in place. ? NEXT = P2** (cross-
platform launcher `indexer --app` + first-run wizard rendered from `/api/capabilities`). P2's
wizard UI is SPA work (khelsutra repo) ? coordinate like P1b.
100    | **P2a** | **ffmpeg resolver helper** wrapping all call sites (`compiler.py`,
`video.py`) + first-run cpu-mock self-test | P0a | No | ? | ? |
101    | P2b | **Cross-platform launcher + install path** (Win/Mac/Linux, no native component):
`uv tool install` / generated script + a pure-Python `indexer --app` (health-poll ?
`webbrowser.open`; free-port fallback; graceful shutdown); per-OS ffmpeg acquisition (winget /
brew / apt) via the P2a resolver; ship `config.json` defaulting truly-local; optional thin per-
OS shortcut | P2a, P1a | No (code-signing only if clickable binaries are ever shipped ? 09) | ?
| ? |
102    | P2c | **In-UI 3-step first-run wizard** (Gemini key / Compute / Videos), rendered from
`/api/capabilities`, en/kn/hi | P1b | No | ? | ? |
103    | P2d | **DB upgrade contract** (backup-before-migrate, downgrade guard, "export my
data") + **uninstall manifest** (keep-my-videos? prompt; scrub `.env` secret) | P0a | No | ? | ?
|
```

525. user (2026-06-21T18:11:09.879Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been
updated successfully. (file state is current in your context ? no need to Read it back)

526. user (2026-06-21T18:11:25.355Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-

indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

527. assistant (2026-06-21T18:11:51.984Z)

Checkpoint reached. Here's the grounded state and what shipped.

Grounding result

- `origin/master` = `ca34787`; ****all 8 of my packaging PRs intact, 0 open PRs****.
- The ****bundled SPA is current**** (khelsutra only moved on `site/`/docs since my pinned `46e0244` ? no `web/` drift).
- ****Coordination note:**** the sibling WIP branch `serving/exposure-safety-preflight` (no PR yet) reworks `startup.py`/`main.py` exposure checks and appears to ****lack my P0c `EXPOSED\_HOST\_ALLOWLIST\_LOOPBACK`**** ? flag for rebase when it opens a PR. I branched every PR off `origin/master` + staged explicit paths, so no collision occurred.

Shipped this session (plan + 8 merged PRs)

- ****Plan:**** docs/PACKAGING_PLAN.md (01/02/bundling decisions locked).
- ****P0 x5**** (#293?#298): app-home, `/api/config` redaction, bind+DNS-rebinding guard, scipy + license-scan & built-wheel CI lanes, single-instance lock.
- ****P1 x2**** (#300, #302): console/wizard endpoints + ****the SPA bundled single-origin into the wheel****.
- ****P2b-launcher**** ([#303](https://github.com/Khelsutra/badminton-highlight-indexer/pull/303)): ****indexer --app**** ? the headline moment. One command ? engine boots ? your browser opens to the bundled UI. Cross-platform, no native component, proven by a subprocess e2e test.

****Net:**** `pip install` the wheel ? `indexer --app` ? a working local app serving the real KhelSutra UI. That's "batteries-included" working end-to-end for the LIGHT tier.

Why I'm pausing the autonomous run here

The next item, ****P2a (the ffmpeg resolver)****, is a ****broad refactor**** ? it routes ****every**** `ffmpeg`/`ffprobe` call site (compiler, audio hit-detector, WASB infer, gemini/yolo segmenters, validators, eval) through one resolver. That carries real regression risk and deserves careful, fresh-context work rather than a rushed pass at the tail of a very long session (8 PRs deep). Stopping on the demonstrable launcher milestone keeps the quality bar intact.

Grounded next steps (in order)

1. ****P2a**** ? ffmpeg resolver (`config-path` ? env ? which ? fail-loud`) + per-OS install offer + first-run cpu-mock self-test. ***(Engine-only; broad but high-value ? makes ffmpeg "just work".)***
2. ****P2d**** ? DB backup-before-migrate + downgrade guard + uninstall manifest. ***(Engine-only.)***
3. ****P2b-install**** ? `uv`/generated cross-platform install script + ship a truly-local default `config.json`.
4. ****P2c**** ? first-run wizard UI (****cross-repo: khelsutra `web/` SPA****, rendered from `/api/capabilities` ? coordinate like P1b).
5. ****P3 (owner-gated)**** ? HEAVY/cloud + the ****real GPU/cloud validation**** where I'd use your VM budget (confirm-cmd-and-cost, delete-after).
6. ****P4 (north star)**** ? ONNX export + training/annotation.

Want me to ****continue with P2a**** (the ffmpeg resolver, with the care it needs) in a fresh push, or pause here for you to review the merged P0/P1/P2-launcher work first?

528. user (2026-06-21T18:19:00.304Z)

yes please continue P2a through P3 where you can spin up GPU enabled GCP VMs and make screencasts to validate the work rigorously

529. assistant (2026-06-21T18:19:55.853Z)

On it ? continuing P2a (ffmpeg resolver), then P2d, then into P3 with real GPU GCP validation + screencasts. I'll announce exact cmd + hourly cost before any cloud spend and delete-after (your standing rule), even though you've authorized it.

Starting ****P2a ? the ffmpeg resolver****. First, re-enumerate every `ffmpeg`/`ffprobe` call site

on current master (verify-before-scoping) so the refactor is complete.

530. user (2026-06-21T18:20:01.029Z)

```
HEAD is now at ca34787 feat(pkg): P2b ? friendly `indexer --app` launcher (boot localhost server
+ open the bundled UI) (#303)
=== ffmpeg/ffprobe as argv[0] in subprocess calls (the sites to wrap) ===
C:/Users/avidu/Projects/bhi-wt-p2a/backend/main.py:338:     ffmpeg = shutil.which("ffmpeg")
C:/Users/avidu/Projects/bhi-wt-p2a/backend/main.py:339:     ffprobe = shutil.which("ffprobe")
C:/Users/avidu/Projects/bhi-wt-p2a/backend/main.py:353:         "ffmpeg": {"ffmpeg":
bool(ffmpeg), "ffprobe": bool(ffprobe),
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/audio/hit_detector.py:34:     cmd =
["ffmpeg", "-y", "-i", video_path, "-vn", "-ac", "1", "-ar", str(sr),
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/segmenters/gemini.py:132:
"ffmpeg", "-y",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/stitcher/compiler.py:113:
"ffprobe", "-v", "error",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/stitcher/compiler.py:249:
"ffmpeg", "-y",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/stitcher/compiler.py:387:
"ffmpeg", "-y",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/validators/format.py:25:
"ffprobe",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/pipeline/validators/pts_continuity.py:17:
"ffprobe",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/tools/rally_slicer.py:109:         ["ffprobe",
"-v", "error", "-show_entries", "format=duration",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/tools/rally_slicer.py:140:         "ffmpeg", "-y",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/utils/video.py:12:         "ffprobe", "-v",
"error", "-show_entries", "format=duration",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/utils/video.py:34:         "ffmpeg", "-y",
C:/Users/avidu/Projects/bhi-wt-p2a/backend/worker/__main__.py:281:         ["ffprobe", "-v",
"error", "-select_streams", "v:0",
```

531. assistant (2026-06-21T18:20:45.958Z)

13 call sites across 9 files (serving path + worker + tools), plus the capabilities probe.
Scoping P2a tightly to the ****resolver + wrapping**** (the cpu-mock self-test belongs with the wizard/first-run ? deferring). Writing the resolver.

532. assistant (2026-06-21T18:20:56.852Z)

"""Single resolver for the ffmpeg / ffprobe binaries (PACKAGING_PLAN.md P2a).

Every ffmpeg/ffprobe invocation in the engine goes through :func:`ffmpeg\_bin` /
:func:`ffprobe\_bin`
so a packaged app can point at a vendored or non-PATH binary via the ``RALLY\_FFMPEG`` /
``RALLY\_FFPROBE`` env vars, there is ONE place to add a config knob or a louder check later, and
a
missing ffmpeg fails with an actionable, OS-specific install hint instead of a cryptic
``FileNotFoundError`` from deep in the pipeline.

DEFAULT (env unset) returns the bare name (``"ffmpeg"`` / ``"ffprobe"``) ? PATH resolution,
exactly
the prior hardcoded literals (byte-identical argv), so this is a pure seam with no behaviour
change.
stdlib-only.
"""

```
from __future__ import annotations
```

```
import os
import platform
import shutil
```

```

class FfmpegNotFound(RuntimeError):
    """Raised by :func:`require_ffmpeg` when ffmpeg/ffprobe can't be located."""

def ffmpeg_bin() -> str:
    """The ffmpeg executable: ``RALLY_FFmpeg`` if set, else ``"ffmpeg"`` (resolved on PATH)."""
    return (os.environ.get("RALLY_FFmpeg") or "").strip() or "ffmpeg"

def ffprobe_bin() -> str:
    """The ffprobe executable: ``RALLY_FFPROBE`` if set, else ``"ffprobe"`` (resolved on
PATH)."""
    return (os.environ.get("RALLY_FFPROBE") or "").strip() or "ffprobe"

def ffmpeg_available() -> bool:
    return shutil.which(ffmpeg_bin()) is not None

def ffprobe_available() -> bool:
    return shutil.which(ffprobe_bin()) is not None

def install_hint() -> str:
    """The OS-appropriate one-line install command (mirrors scripts/doctor.py)."""
    system = platform.system()
    if system == "Windows":
        return "winget install Gyan.FFmpeg"
    if system == "Darwin":
        return "brew install ffmpeg"
    return "sudo apt-get install -y ffmpeg (or your distro's package manager)"

def require_ffmpeg() -> None:
    """Raise :class:`FfmpegNotFound` (with an OS-specific install hint) if ffmpeg/ffprobe is
missing.

    A loud, actionable pre-flight for entry points that are about to shell out ? used by the
    first-run self-test / wizard; the per-call-site argv still goes through
:func:`ffmpeg_bin`.
    """
    missing = [name for name, ok in (("ffmpeg", ffmpeg_available()), ("ffprobe",
ffprobe_available())) if not ok]
    if missing:
        raise FfmpegNotFound(
            f"{' and '.join(missing)} not found (PATH or RALLY_FFmpeg/RALLY_FFPROBE). FFmpeg is
"
            f"required for video processing. Install it: {install_hint()}")

```

533. user (2026-06-21T18:20:56.870Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p2a\backend\utils\ffmpeg.py (file state is current in your context ? no need to Read it back)

534. assistant (2026-06-21T18:21:09.338Z)

Now wrap the call sites. Let me read the smaller files fully and the relevant sections of the larger ones.

535. user (2026-06-21T18:21:09.654Z)

```

1     import subprocess
2     import logging
3     from typing import Optional
4

```

```

5     logger = logging.getLogger(__name__)
6
7     def get_video_duration(video_path: str) -> float:
8         """Returns the exact duration of a video file in seconds using ffprobe.
9         Returns 0.0 on failure."""
10        try:
11            duration_output = subprocess.check_output([
12                "ffprobe", "-v", "error", "-show_entries", "format=duration",
13                "-of", "default=noprint_wrappers=1:nokey=1", video_path
14            ]).decode().strip()
15            return float(duration_output)
16        except Exception as e:
17            logger.error(f"Could not detect video duration for {video_path}: {e}")
18            return 0.0
19
20    def extract_video_chunk(video_path: str, start_time: float, duration: float,
21                           output_path: str, reencode: bool = False,
22                           scale_width: Optional[int] = None) -> bool:
23        """Extracts a chunk from a video file.
24
25        Args:
26            reencode: If True, re-encodes with libx264 (slower but frame-accurate).
27                     If False, uses stream-copy (fast but cuts at nearest keyframe).
28                     Use True for short clips where keyframe drift matters.
29            scale_width: If set (requires reencode), downscale to this width (height auto,
30                       even). Useful for high-res sources headed to an API that analyses at
31                       low res anyway ? shrinks uploads with no analysis loss.
32
33        """
34        cmd = [
35            "ffmpeg", "-y",
36            "-ss", str(start_time),
37            "-t", str(duration),
38            "-i", video_path,
39        ]
40        if reencode:
41            if scale_width:
42                cmd += ["-vf", f"scale={int(scale_width)}:-2"]
43            cmd += ["-c:v", "libx264", "-preset", "ultrafast", "-crf", "23", "-c:a", "aac"]
44        else:
45            cmd += ["-c", "copy"]
46        cmd.append(output_path)
47        try:
48            subprocess.run(cmd, check=True, stdout=subprocess.DEVNULL,
49                           stderr=subprocess.DEVNULL)
50            return True
51        except Exception as e:
52            logger.error(f"Failed to extract chunk {start_time}s - {start_time+duration}s: {e}")
53            return False

```

536. user (2026-06-21T18:21:14.617Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\utils\video.py has been updated successfully. (file state is current in your context ? no need to Read it back)

537. user (2026-06-21T18:21:19.388Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\utils\video.py has been updated successfully. (file state is current in your context ? no need to Read it back)

538. user (2026-06-21T18:21:23.961Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\utils\video.py has been updated

successfully. (file state is current in your context ? no need to Read it back)

539. user (2026-06-21T18:21:31.217Z)

```
1     import subprocess
2     import json
3     import os
4     from backend.pipeline.validators.base import VideoValidator, ValidationResult
5     from typing import Dict, Any
6
7     class FormatValidator(VideoValidator):
8         @property
9         def name(self) -> str:
10             return "format_check"
11
12         @property
13         def description(self) -> str:
14             return "Verifies the video is a valid, uncorrupted MP4 container and extracts
metadata."
15
16         def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
17             if not os.path.exists(video_path):
18                 return ValidationResult(
19                     validator_name=self.name,
20                     passed=False,
21                     message=f"File not found: {video_path}"
22                 )
23
24             cmd = [
25                 "ffprobe",
26                 "-v", "error",
27                 "-select_streams", "v:0",
28                 "-show_format",
29                 "-show_streams",
30                 "-of", "json",
```

540. user (2026-06-21T18:21:31.993Z)

```
1     import subprocess
2     from backend.pipeline.validators.base import VideoValidator, ValidationResult
3     from typing import Dict, Any
4
5     class PTSContinuityValidator(VideoValidator):
6         @property
7         def name(self) -> str:
8             return "pts_continuity_check"
9
10        @property
11        def description(self) -> str:
12            return "Analyzes Presentation Time Stamps (PTS) of keyframes to check for
timeline gaps, jumps, or editing splices."
13
14        def validate(self, video_path: str, config: Dict[str, Any], context: Dict[str, Any])
-> ValidationResult:
15            # Command to extract packet PTS and flags (extremely fast, works on H.264 &
H.265)
16            cmd = [
17                "ffprobe",
18                "-v", "error",
19                "-select_streams", "v:0",
20                "-show_entries", "packet=pts_time,flags",
21                "-of", "csv=p=0",
22                video_path
```

541. user (2026-06-21T18:21:42.031Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\validators\format.py has been updated successfully. (file state is current in your context ? no need to Read it back)

542. user (2026-06-21T18:21:47.297Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\validators\format.py has been updated successfully. (file state is current in your context ? no need to Read it back)

543. user (2026-06-21T18:21:53.025Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\validators\pts_continuity.py has been updated successfully. (file state is current in your context ? no need to Read it back)

544. user (2026-06-21T18:22:05.133Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\validators\pts_continuity.py has been updated successfully. (file state is current in your context ? no need to Read it back)

545. user (2026-06-21T18:22:11.241Z)

```
1      """Audio shuttle-hit ("thwack") detection ? local, dependency-light (numpy + ffmpeg).
2
3      Stages 0?1 of docs/AUDIO_RALLY_DETECTION_PLAN.md: extract the audio track, build a
4      spectral-flux onset envelope (with pre-emphasis to boost the broadband HF transient
5      a racket?shuttle contact makes), threshold it adaptively, and return (timestamp,
6      strength) hit candidates. Pure numpy + stdlib `wave` (NO scipy); ffmpeg (subprocess)
7      only for extraction. Consumed by the E1 go/no-go probe; later feeds rally fusion
8      (ADR-007) if E1 passes. WASB is never modified ? this is a separate signal.
9      """
10     from __future__ import annotations
11
12     import os
13     import wave
14     import tempfile
15     import subprocess
16     from dataclasses import dataclass
17     from typing import List, Optional, Tuple
18
19     import numpy as np
20
21
22     @dataclass
23     class HitEvent:
24         t: float          # seconds
25         strength: float    # onset-envelope peak value (arbitrary units)
26
27
28     def extract_audio(video_path: str, out_wav: Optional[str] = None, sr: int = 16000) ->
29     str:
30         """ffmpeg ? mono 16-bit PCM WAV at `sr`. Returns the wav path (temp if not given).
31         Audio-only decode, so it's fast even for multi-GB video."""
32         if out_wav is None:
33             fd, out_wav = tempfile.mkstemp(suffix=".wav")
34             os.close(fd)
35             cmd = ["ffmpeg", "-y", "-i", video_path, "-vn", "-ac", "1", "-ar", str(sr),
36                 "-c:a", "pcm_s16le", out_wav]
37             subprocess.run(cmd, check=True, stdout=subprocess.DEVNULL,
38                 stderr=subprocess.DEVNULL)
39             return out_wav
40
41     def load_wav(path: str) -> Tuple[np.ndarray, int]:
```

546. user (2026-06-21T18:22:16.120Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\audio\hit_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

547. user (2026-06-21T18:22:21.298Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\audio\hit_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

548. user (2026-06-21T18:22:25.640Z)

```
125         if os.path.exists(temp_trimmed_path):
126             try:
127                 os.remove(temp_trimmed_path)
128             except Exception:
129                 pass
130
131         cmd = [
132             "ffmpeg", "-y",
133             "-ss", "0",
134             "-t", str(debug_duration),
135             "-i", video_path,
136             "-c", "copy",
```

549. user (2026-06-21T18:22:30.566Z)

```
1:import os
2:import time
3:import logging
4:import subprocess
5:import concurrent.futures
6:from typing import List, Dict, Any, Tuple, Callable
7:from backend.pipeline.segmenters.base import BasePlaySegmenter, segmenter_registry
8:from backend.utils.video import get_video_duration, extract_video_chunk
9:from backend.utils.paths import output_base
10:from backend.sports import sport_registry
11:from backend.providers import provider_registry
12:from backend.reporting import run_signals
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

550. user (2026-06-21T18:22:35.517Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\segmenters\gemini.py has been updated successfully. (file state is current in your context ? no need to Read it back)

551. user (2026-06-21T18:22:42.446Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\segmenters\gemini.py has been updated successfully. (file state is current in your context ? no need to Read it back)

552. user (2026-06-21T18:22:48.932Z)

```
1:import os
2:import subprocess
3:import tempfile
4:import logging
--- sites ---
    [
        "ffprobe", "-v", "error",
        "-show_entries", "format=duration",
base_cmd = [
    "ffmpeg", "-y",
    "-ss", str(round(padded_start, 3)),
```

```
concat_cmd = [  
    "ffmpeg", "-y",  
    "-f", "concat",
```

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

553. user (2026-06-21T18:22:53.918Z)

```
1     import os  
2     import subprocess  
3     import tempfile  
4     import logging  
5     from pathlib import Path  
6     from typing import Any, Dict, List, Optional, Tuple, Union  
7  
8     logger = logging.getLogger(__name__)  
9  
10    # Bundled, commercially-clean default watermark font (Apache-2.0). Shipped so the  
overlay  
11    # renders out-of-the-box without a system fontconfig setup (a common gap on Windows,  
where  
12    # `font='Sans'` fails). Overridable via stitching.watermark.font_path. This module lives  
at
```

554. user (2026-06-21T18:23:01.657Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

555. user (2026-06-21T18:23:06.489Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

556. user (2026-06-21T18:23:11.855Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

557. user (2026-06-21T18:23:16.859Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\pipeline\stitcher\compiler.py has been updated successfully. (file state is current in your context ? no need to Read it back)

558. assistant (2026-06-21T18:23:22.879Z)

Now rally_slicer.py, worker, and the main.py capabilities probe. Let me read those sites.

559. user (2026-06-21T18:23:25.668Z)

```
=== rally_slicer imports + sites ===  
20:import argparse  
21:import csv  
22:import os  
23:import subprocess  
27:from backend.eval.annotations import parse_timestamp  
    out = subprocess.run(  
        ["ffprobe", "-v", "error", "-show_entries", "format=duration",  
        "-of", "default=noprint_wrappers=1:nokey=1", video_path],  
    cmd = [  
        "ffmpeg", "-y",  
        "-ss", str(w.start), "-to", str(w.end), "-i", video_path,  
=== worker import area + site ===  
18:import argparse  
19:import csv
```

```

20:import json
21:import logging
22:import os
23:import sys
24:import tempfile
27:from backend import storage
28:from backend.compute import ComputeJobSpec, get_compute_backend
29:from backend.config import (
    out = subprocess.check_output(
        ["ffprobe", "-v", "error", "-select_streams", "v:0",
        "-show_entries", "stream=r_frame_rate,width", "-of", "json", path],
=== main.py capabilities probe ===
    indexing = getattr(global_config, "indexing", None)
    deployment = getattr(global_config, "deployment", None)
    ffmpeg = shutil.which("ffmpeg")
    ffprobe = shutil.which("ffprobe")
    # WASB weights are best-effort: env first (matches native_wasb_runner), then config.
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

560. user (2026-06-21T18:23:30.124Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

561. user (2026-06-21T18:23:33.942Z)

```

20     import argparse
21     import csv
22     import os
23     import subprocess
24     from dataclasses import dataclass
25     from typing import Dict, List, Optional, Tuple
26
27     from backend.eval.annotations import parse_timestamp
28
29     # Match VideoCompiler defaults so sliced clips frame a rally the same way.

```

562. user (2026-06-21T18:23:39.630Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\tools\rally_slicer.py has been updated successfully. (file state is current in your context ? no need to Read it back)

563. user (2026-06-21T18:23:46.703Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\tools\rally_slicer.py has been updated successfully. (file state is current in your context ? no need to Read it back)

564. user (2026-06-21T18:23:52.352Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\tools\rally_slicer.py has been updated successfully. (file state is current in your context ? no need to Read it back)

565. user (2026-06-21T18:23:57.986Z)

```

27     from backend import storage
28     from backend.compute import ComputeJobSpec, get_compute_backend
29     from backend.config import (
30         StartupCheckError,
31         enforce_startup_checks,
32         load_config,
33         probe_cuda_available,
34     )
35     from backend.infrastructure.database import Database
36     from backend.pipeline.segmenters.base import segmenter_registry

```


566. user (2026-06-21T18:24:02.464Z)

```
36     from backend.pipeline.segmenters.base import segmenter_registry
37     from backend.pipeline.validators.base import registry as validator_registry
38     from backend.reporting import emit_indexer_report
39     from backend.results import Failure
40     from backend.utils.hashing import compute_video_id
41
42     logger = logging.getLogger("backend.worker")
43
```

567. user (2026-06-21T18:24:07.851Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

568. user (2026-06-21T18:24:13.191Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\worker__main__.py has been updated successfully. (file state is current in your context ? no need to Read it back)

569. assistant (2026-06-21T18:24:20.356Z)

Finally route the main.py capabilities probe through the resolver (so the env override reflects).

570. user (2026-06-21T18:24:22.272Z)

<tool_use_error>File has not been read yet. Read it first before writing to it.</tool_use_error>

571. user (2026-06-21T18:24:27.776Z)

```
44     from backend.utils.paths import resolve_under_root, safe_output_filename,
validate_input_path
45     from backend.utils.redact import redact_secrets # P0b: mask secret-shaped fields in
/api/config
46     from backend.utils.single_instance import acquire_lock # P0e: one instance per database
```

572. user (2026-06-21T18:24:33.768Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

573. user (2026-06-21T18:24:39.178Z)

The file C:\Users\avidu\Projects\bhi-wt-p2a\backend\main.py has been updated successfully. (file state is current in your context ? no need to Read it back)

574. assistant (2026-06-21T18:24:50.338Z)

All 13 sites wrapped + the capabilities probe. Now the resolver tests.

575. assistant (2026-06-21T18:25:01.794Z)

"""The single ffmpeg/ffprobe resolver (PACKAGING_PLAN.md P2a)."""

```
import subprocess
```

```
import pytest
```

```
from backend.utils import ffmpeg as ff
```

```

@pytest.fixture(autouse=True)
def _clear_ffmpeg_env(monkeypatch):
    monkeypatch.delenv("RALLY_FFMPEG", raising=False)
    monkeypatch.delenv("RALLY_FFPROBE", raising=False)

```

---- default (env unset) is byte-identical to the old hardcoded literals ----

```

def test_default_bins_are_bare_names():
    assert ff.ffmpeg_bin() == "ffmpeg"
    assert ff.ffprobe_bin() == "ffprobe"

def test_env_override(monkeypatch):
    monkeypatch.setenv("RALLY_FFMPEG", "/opt/ff/ffmpeg")
    monkeypatch.setenv("RALLY_FFPROBE", "/opt/ff/ffprobe")
    assert ff.ffmpeg_bin() == "/opt/ff/ffmpeg"
    assert ff.ffprobe_bin() == "/opt/ff/ffprobe"

def test_blank_env_falls_back_to_default(monkeypatch):
    monkeypatch.setenv("RALLY_FFMPEG", " ")
    assert ff.ffmpeg_bin() == "ffmpeg"

```

---- availability + require ----

```

def test_available_reflects_which(monkeypatch):
    monkeypatch.setattr(ff.shutil, "which", lambda name: "/usr/bin/" + name)
    assert ff.ffmpeg_available() is True and ff.ffprobe_available() is True
    monkeypatch.setattr(ff.shutil, "which", lambda name: None)
    assert ff.ffmpeg_available() is False and ff.ffprobe_available() is False

def test_require_ffmpeg_noop_when_present(monkeypatch):
    monkeypatch.setattr(ff.shutil, "which", lambda name: "/usr/bin/" + name)
    ff.require_ffmpeg() # must not raise

def test_require_ffmpeg_raises_with_hint_when_missing(monkeypatch):
    monkeypatch.setattr(ff.shutil, "which", lambda name: None)
    with pytest.raises(ff.FfmpegNotFound) as exc:
        ff.require_ffmpeg()
    msg = str(exc.value)
    assert "ffmpeg" in msg and "ffprobe" in msg
    assert ff.install_hint() in msg # actionable, OS-specific

@pytest.mark.parametrize("system,expect", [
    ("Windows", "winget"), ("Darwin", "brew"), ("Linux", "apt-get")])
def test_install_hint_per_os(monkeypatch, system, expect):
    monkeypatch.setattr(ff.platform, "system", lambda: system)
    assert expect in ff.install_hint()

```

---- a wrapped call site actually uses the resolved binary ----

```

def test_call_site_uses_resolved_binary(monkeypatch):
    from backend.utils import video
    monkeypatch.setenv("RALLY_FFPROBE", "/custom/ffprobe")
    captured = {}

    def _fake_check_output(argv, *a, **k):

```

```
captured["argv"] = argv
return b"12.5"
```

```
monkeypatch.setattr(subprocess, "check_output", _fake_check_output)
assert video.get_video_duration("x.mp4") == 12.5
assert captured["argv"][0] == "/custom/ffprobe" # the override flows through to argv[0]
```

576. user (2026-06-21T18:25:01.914Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p2a\tests\test_ffmpeg_resolver.py
(file state is current in your context ? no need to Read it back)

577. user (2026-06-21T18:25:11.777Z)

```
..... [100%]
10 passed in 0.18s
=== ruff + mypy (changed) ===
All checks passed!
Success: no issues found in 1 source file
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

578. assistant (2026-06-21T18:25:21.717Z)

10 tests pass, ruff/mypy clean. Now the full suite ? the key check that the wrapping is byte-identical (env unset) and broke no subprocess-mocking test.

579. user (2026-06-21T18:27:49.095Z)

```
backend\utils\ffmpeg.py          24      0   100%
TOTAL                          9912  1475   85%
Required test coverage of 70% reached. Total coverage: 85.12%
1312 passed, 7 skipped, 7 warnings in 142.36s (0:02:22)
RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

580. assistant (2026-06-21T18:27:58.110Z)

Full suite green (1312), `ffmpeg.py` 100%, byte-identical (no test broke). Let me confirm no `ffmpeg`/`ffprobe` argv literal was missed.

581. user (2026-06-21T18:28:01.205Z)

```
=== remaining unwrapped argv literals (should be NONE in backend/, excluding the resolver +
hints) ===
(empty above = all wrapped)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

582. assistant (2026-06-21T18:28:14.223Z)

Complete ? every argv literal wrapped, byte-identical, full suite green, resolver 100%. (This is a mechanical byte-neutral refactor verified by the completeness grep + a call-site wiring test + 1312 green, so I'll commit directly rather than spend a review pass ? conserving budget for the P3 GPU work.) Committing P2a.

583. user (2026-06-21T18:28:30.293Z)

```
origin/master: e99d0c3 feat(stitch): localize the reel watermark to the UI locale (+ OFL Noto
fonts) (#304)
warning: in the working copy of 'backend/utils/ffmpeg.py', LF will be replaced by CRLF the next
time Git touches it
warning: in the working copy of 'tests/test_ffmpeg_resolver.py', LF will be replaced by CRLF the
next time Git touches it
=== staged ===
```

```
11
7ad0f1c feat(pkg): P2a ? single ffmpeg/ffprobe resolver (env-overridable; one fail-loud place)
(default-OFF)
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch]      feat/packaging-p2a-ffmpeg -> feat/packaging-p2a-ffmpeg
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

584. assistant (2026-06-21T18:28:54.167Z)

origin/master moved to `e99d0c3` (#304 ? **reel watermark localization, which touches `compiler.py` ? a file my P2a also edits). Potential conflict. Let me check the overlap and rebase if needed before opening the PR.

585. user (2026-06-21T18:28:57.033Z)

```
=== what #304 changed in compiler.py (vs my base ca34787) ===
```

```
=== attempt rebase onto current origin/master ===
Rebasing (1/1)
```

Successfully rebased and updated refs/heads/feat/packaging-p2a-ffmpeg.
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

586. assistant (2026-06-21T18:29:13.979Z)

Rebase succeeded **cleanly** ? #304 didn't overlap my ffmpeg sites (it added `watermark\_i18n`). My P2a now sits on top of `e99d0c3`. I must re-verify post-rebase (the tree changed) before pushing.

587. user (2026-06-21T18:31:02.409Z)

```
=== HEAD now (1 commit on e99d0c3) ===
8f915cc feat(pkg): P2a ? single ffmpeg/ffprobe resolver (env-overridable; one fail-loud place)
(default-OFF)
=== re-verify post-rebase: ruff + full suite ===
All checks passed!
1319 passed, 7 skipped in 102.02s (0:01:42)
=====
RESULT: PASS - all tests green
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

588. assistant (2026-06-21T18:31:08.220Z)

Post-rebase clean (ruff clean, 1319 passed ? #304's watermark tests included). Force-pushing the rebased branch + opening the PR.

589. user (2026-06-21T18:31:23.566Z)

```
To https://github.com/Khelsutra/badminton-highlight-indexer.git
+ 7ad0f1c...8f915cc feat/packaging-p2a-ffmpeg -> feat/packaging-p2a-ffmpeg (forced update)
https://github.com/Khelsutra/badminton-highlight-indexer/pull/305
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

590. assistant (2026-06-21T18:31:29.854Z)

PR #305 opened: [feat(pkg): P2a ? ffmpeg resolver](https://github.com/Khelsutra/badminton-highlight-indexer/pull/305) (rebased onto current master). Watching CI.

591. user (2026-06-21T18:34:16.241Z)

```
Syntax sweep + import smoke (no GPU stack)      pass      23s
https://github.com/Khelsutra/badminton-highlight-
```

```
indexer/actions/runs/27913607580/job/82594888473
Built wheel smoke (install + launch) pass 1m15s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27913607580/job/82594888495
Dependency license scan (no copyleft in runtime deps) pass 35s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913607580/job/82594888484
Full test suite (offline, mocked) pass 2m35s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27913607580/job/82594888508
Golden fixtures cross-OS (macos-latest) pass 17s https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27913607580/job/82594888490
Golden fixtures cross-OS (ubuntu-latest) pass 25s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913607580/job/82594888500
Golden fixtures cross-OS (windows-latest) pass 1m12s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913607580/job/82594888549
Mypy (lenient baseline) pass 31s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913607580/job/82594888516
Ruff lint pass 11s https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913607580/job/82594888468
Syntax sweep + import smoke (no GPU stack) pass 23s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27913607580/job/82594888473
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

592. assistant (2026-06-21T18:34:19.477Z)

All 9 lanes green. Merging P2a and cleaning up.

593. user (2026-06-21T18:34:28.406Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
e99d0c3..fa7ee3f master -> origin/master
=== master ===
fa7ee3f feat(pkg): P2a ? single ffmpeg/ffprobe resolver (env-overridable; one fail-loud place)
(default-OFF) (#305)
=== PR #305 ===
MERGED
```

594. assistant (2026-06-21T18:36:11.902Z)

P2a ****merged**** (master `fa7ee3f`) ? that's ****9 PRs this session**** (plan + P0×5 + P1×2 +
P2-launcher + P2a). Updating the tracker, then grounding the GPU/P3 validation.

595. user (2026-06-21T18:36:16.443Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been
updated successfully. (file state is current in your context ? no need to Read it back)

596. assistant (2026-06-21T18:36:21.995Z)

Now grounding the GPU/P3 validation: GCP auth/quota, any running VMs (cost audit), the provision
scripts, and the GCS corpus (smoke video + weights).

597. user (2026-06-21T18:36:43.007Z)

```
=== gcloud auth + project ===
avi.dullu@gmail.com
gen-lang-client-0306026826

=== any running VMs (cost audit) ===
Listed 0 items.
```

```
=== GPU quota (GPUS_ALL_REGIONS) ===
ERROR: (gcloud.compute.regions.describe) Transform function expected [default
value(quotas.filter(metric=GPUS_ALL_REGIONS).limit *HERE* )].
```

```
=== deploy/gcp scripts present ===
```

```
README.md
create_gpu_vm.sh
ops-agent-startup.sh
provision.sh
teardown.sh
```

```
=== GCS corpus: smoke video + weights present? ===
```

```
gs://khelsutra-rally-corpus/highlights/
gs://khelsutra-rally-corpus/labels/
gs://khelsutra-rally-corpus/models/
gs://khelsutra-rally-corpus/outputs/
gs://khelsutra-rally-corpus/promo/
gs://khelsutra-rally-corpus/videos/
gs://khelsutra-rally-corpus/weights/
```

598. assistant (2026-06-21T18:37:21.319Z)

GCP is **verified ready**: authed (`avi.dullu@gmail.com` / `gen-lang-client-0306026826`), **0** running VMs (\$0 now), `deploy/gcp/{create\_gpu\_vm,provision,teardown,ops-agent-startup}.sh` present, and the corpus (`videos/`, `models/`, `weights/`) staged. Let me record the validation plan in memory, then give you a grounded recommendation.

599. user (2026-06-21T18:37:32.090Z)

The file C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

600. assistant (2026-06-21T18:37:48.962Z)

Here's the honest, grounded recommendation.

Done: P2a merged ? 9 PRs this session

[#305](https://github.com/Khelsutra/badminton-highlight-indexer/pull/305) ffmpeg/ffprobe resolver (env-overridable, fail-loud, all 13 call sites wrapped, byte-identical). Master is `fa7ee3f`; the full P0+P1+P2-launcher+P2a stack is live and the LIGHT-tier app is functionally proven end-to-end (the `indexer --app` subprocess test boots the wheel and serves the bundled SPA).

GPU/P3 validation ? verified ready, and here's my call on ***when***

The infra checks out for the rigorous GPU validation you want:

- **Cost**: L4 `g2-standard-4` (asia-south1) ? **\$0.70/hr**; a 30?60 min validation ? **\$0.35?0.70** ? trivially within the \$100 cap.
- **Plan**: `deploy/gcp/create\_gpu\_vm.sh` ? install **the wheel** on the fresh GPU box + the WASB env ? validate the ***new packaging*** end-to-end on the real target (wheel ? app-home ? `indexer --app` serving the bundled SPA ? ffmpeg resolver) **and** native-WASB on the staged smoke clip ? screencast ? `teardown.sh` (defaults to ***delete***) + confirm \$0.

My recommendation: run this as a focused, fresh push rather than now ? for two concrete reasons, not reluctance:

1. **Billable-resource safety**. A GPU VM must be torn down start-to-finish. I'm 9 PRs deep in this session; if I start a billable VM and lose context mid-operation, it keeps billing ? exactly the mistake [[gcp-vm-always-delete]] exists to prevent (you've been burned twice). A GPU run deserves undistracted attention with teardown as priority #1.
2. **A known rabbit hole**. Last session, native-WASB on a ***fresh*** build returned **0** rallies ? reproducing the 22-rally result needs the engine session's exact tuned `worker infer detector\_impl=native` config (logged as "the missing piece"). So the HEAVY-tier validation isn't a clean 30-min win; it may need real debugging time. I don't want to half-do it at the tail of a

long session.

I've recorded the full ready-to-execute plan + cost + that blocker in memory, so a fresh push can start the VM and run it cleanly.

Your call

601. user (2026-06-21T19:01:04.539Z)

Your questions have been answered: "The GPU/P3 validation is verified-ready (~\$0.50 run, well under the cap). How do you want to handle the billable GPU VM run given this session is already 9 PRs deep?"="1. with the remaining items (in 3) as a continuation with the same repo diligence.". You can now continue with these answers in mind.

602. assistant (2026-06-21T19:01:33.892Z)

Got it ? ****GPU/P3 deferred to a fresh focused push**** (responsible for the billable VM), and I'll ****continue now with the remaining engine-only P2 items**** (P2d DB-upgrade safety, P2b-install, local LIGHT-tier screencast) with the same diligence.

Starting ****P2d ? DB-upgrade safety**** (backup-before-migrate + downgrade guard) ? table-stakes for an auto-updating downloadable app. Worktree + reading the migration ladder.

603. user (2026-06-21T19:01:37.112Z)

```
HEAD is now at f01dc85 feat(db): v7 ? record the UI locale per job (PMF analytics) +
count_jobs_by_locale (#306)
=== migration ladder (version read/dispatch/set) ===
        logger.error(f"{error_msg}: {e}")
        return False
```

```
@staticmethod
```

```
def _add_column_idempotent(cursor: sqlite3.Cursor, ddl: str) -> None:
    """Apply a bare ``ALTER TABLE ... ADD COLUMN`` migration step re-runnably.
```

```

    A DDL ``ALTER`` auto-commits independently of the surrounding ``with conn``
    transaction, so an interrupted v3/v4 block can leave the column added while
    ``user_version`` stays un-bumped ? and the next open would re-run the same
    ``ALTER`` and raise ``OperationalError: duplicate column name``, bricking every
    future ``Database(...)`` construction (``_init_db`` runs in ``__init__``). Swallow
    ONLY that specific error so the ladder stays crash-safe, matching the re-runnable
    ``CREATE ... IF NOT EXISTS`` v1/v2/v5 blocks.
    """
```

```
    try:
        cursor.execute(ddl)
    except sqlite3.OperationalError as exc:
        if "duplicate column name" not in str(exc).lower():
            raise
```

```
def _init_db(self):
    """Initializes tables for videos, play_segments, and events."""
    with self._connect() as conn:
        cursor = conn.cursor()
```

```
        self._create_base_tables(cursor)
```

```
        # Versioned migrations live here. PRAGMA user_version is the schema
        # contract ? bump it for every change and add an ordered `if version < N`
        # block below. Tests assert the expected version, so accidental drift fails CI.
        version = cursor.execute("PRAGMA user_version").fetchone()[0]
```

```
        if version < 1:
            self._migrate_to_v1(cursor)
```

```

if version < 2:
    self._migrate_to_v2(cursor)

if version < 3:
    self._migrate_to_v3(cursor)

if version < 4:
    self._migrate_to_v4(cursor)

if version < 5:
    self._migrate_to_v5(cursor)

if version < 6:
    self._migrate_to_v6(cursor)

if version < 7:
    self._migrate_to_v7(cursor)
# future: if version < 8: cursor.execute("ALTER TABLE ..."); PRAGMA user_version = 8

conn.commit()

```

604. assistant (2026-06-21T19:02:13.908Z)

Master moved to `f01dc85` (#306 added `**v7**` ? UI-locale per job), so the schema max is now `**7**`. P2d must derive/match that. Let me read `\_connect` (WAL) + the v7 tail + the existing version-assertion test.

605. user (2026-06-21T19:02:17.499Z)

=== _connect (WAL?) + db_path ===

```

class Database:
    def __init__(self, db_path: str):
        self.db_path = db_path
        self._init_db()

    @contextmanager
    def _connect(self):
        """Transaction scope that also CLOSES the connection.

        ``with sqlite3.connect(...)`` alone is only a commit/rollback scope ? it
        never closes, so every call leaked its connection to GC (lingering file
        handles on the WAL sidecars under Windows). Internal methods use this.
        """
        conn = self._get_connection()
        try:
            with conn:
                yield conn
        finally:
            conn.close()

    def _get_connection(self):
        # busy_timeout: wait (not fail) on a locked DB ? needed once the async job model
        # adds concurrent writers (a swallowed 'database is locked' silently drops segments).
        conn = sqlite3.connect(self.db_path, timeout=30.0)
        conn.row_factory = sqlite3.Row
        # Enforce ON DELETE CASCADE / SET NULL ? SQLite leaves FKs off per-connection.
        conn.execute("PRAGMA foreign_keys = ON")
        # WAL allows concurrent readers during a write; busy_timeout bounds lock waits.
        conn.execute("PRAGMA busy_timeout = 30000")
        try:
            conn.execute("PRAGMA journal_mode = WAL")
        except sqlite3.OperationalError:
            pass # e.g. a read-only/networked FS that can't do WAL ? degrade gracefully
        return conn

```



```

=== v7 migration sets user_version=7? ===
1:import sqlite3
2:import json
3:import logging
342:     def _migrate_to_v7(self, cursor):
349:         cursor.execute("PRAGMA user_version = 7")
=== tests asserting user_version ===
tests/test_async_jobs.py:10:- migration: fresh DB lands at user_version 2 with a jobs table; v1
DBs upgrade in place
tests/test_async_jobs.py:83:         assert conn.execute("PRAGMA user_version").fetchone()[0] ==
7
tests/test_async_jobs.py:101:         conn.execute("PRAGMA user_version = 1")
tests/test_async_jobs.py:107:         assert c.execute("PRAGMA user_version").fetchone()[0] == 7
tests/test_billing.py:39:         ver = conn.execute("PRAGMA user_version").fetchone()[0]
tests/test_database.py:82: def test_schema_version_is_current(tmp_path):
tests/test_database.py:86:     version = conn.execute("PRAGMA user_version").fetchone()[0]
tests/test_database.py:89:     assert version == 7
tests/test_database.py:108:     version = conn.execute("PRAGMA user_version").fetchone()[0]
tests/test_database.py:114:     assert version == 7          # v7 = job UI locale (PMF analytics); v6
= cost ledger (M8)
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

606. assistant (2026-06-21T19:02:44.865Z)

Clear: schema max = **7**, `\_init\_db` runs in `\_\_init\_\_` under WAL, and there are existing tests (a v1?reopen upgrade test + a `== 7` pin). P2d design: `SCHEMA\_VERSION` constant + downgrade guard (refuse a newer-than-known DB) + backup-before-migrate (checkpoint + copy) + a `backup\_database()` helper. Let me read the imports + insertion point.

607. user (2026-06-21T19:02:45.665Z)

```

1     import sqlite3
2     import json
3     import logging
4     from contextlib import contextmanager
5     from typing import List, Dict, Any, Optional, Tuple
6
7     from backend import billing # pure cost math (no backend imports ? no cycle)
8
9     logger = logging.getLogger(__name__)

```

608. user (2026-06-21T19:02:55.538Z)

The file C:\Users\avidu\Projects\bhi-wt-p2d\backend\infrastructure\database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

609. assistant (2026-06-21T19:02:59.037Z)

Now insert the downgrade guard + backup-before-migrate in `\_init\_db`.

610. user (2026-06-21T19:03:05.531Z)

The file C:\Users\avidu\Projects\bhi-wt-p2d\backend\infrastructure\database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

611. assistant (2026-06-21T19:03:12.378Z)

Now add the backup helpers. Let me find the end of `\_init\_db`.

612. user (2026-06-21T19:03:13.195Z)

```

126         if version < 3:
127             self._migrate_to_v3(cursor)
128

```

```

129             if version < 4:
130                 self._migrate_to_v4(cursor)
131
132             if version < 5:
133                 self._migrate_to_v5(cursor)
134
135             if version < 6:
136                 self._migrate_to_v6(cursor)
137
138             if version < 7:
139                 self._migrate_to_v7(cursor)

```

613. user (2026-06-21T19:03:17.115Z)

```

139             self._migrate_to_v7(cursor)
140             # future: if version < 8: cursor.execute("ALTER TABLE ..."); PRAGMA
user_version = 8
141
142             conn.commit()
143
144             def _create_base_tables(self, cursor):
145                 """Create the videos / play_segments / events base tables (idempotent)."""
146                 # Videos Table
147                 cursor.execute("""
148                     CREATE TABLE IF NOT EXISTS videos (

```

614. user (2026-06-21T19:03:31.000Z)

The file C:\Users\avidu\Projects\bhi-wt-p2d\backend\infrastructure\database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

615. assistant (2026-06-21T19:04:07.606Z)

Now the tests for the upgrade-safety contract.

616. assistant (2026-06-21T19:04:18.844Z)

```

"""DB upgrade-safety contract for an auto-updating downloadable app (PACKAGING_PLAN.md P2d)."""

import os
import sqlite3

import pytest

from backend.infrastructure.database import SCHEMA_VERSION, Database, SchemaTooNewError

def _user_version(path: str) -> int:
    conn = sqlite3.connect(path)
    try:
        return int(conn.execute("PRAGMA user_version").fetchone()[0])
    finally:
        conn.close()

def _set_user_version(path: str, v: int) -> None:
    conn = sqlite3.connect(path)
    try:
        conn.execute(f"PRAGMA user_version = {int(v)}")
        conn.commit()
    finally:
        conn.close()

```

```

def test_schema_version_matches_ladder(tmp_path):
    """Ties the SCHEMA_VERSION constant to the migration ladder: a freshly-built DB lands
    exactly at
    SCHEMA_VERSION. Adding a _migrate_to_vN without bumping the constant fails here."""
    p = str(tmp_path / "t.db")
    Database(p)
    assert _user_version(p) == SCHEMA_VERSION

def test_fresh_db_makes_no_backup(tmp_path):
    Database(str(tmp_path / "t.db"))
    assert not any(".bak" in f.name for f in tmp_path.iterdir()) # version 0 ? nothing to back
    up

def test_reopen_current_db_no_backup_no_error(tmp_path):
    p = str(tmp_path / "t.db")
    Database(p) # ? SCHEMA_VERSION
    Database(p) # reopen at current version ? no migration runs
    assert not any(".bak" in f.name for f in tmp_path.iterdir())

def test_downgrade_guard_refuses_newer_db(tmp_path):
    p = str(tmp_path / "t.db")
    Database(p)
    _set_user_version(p, SCHEMA_VERSION + 5) # a DB from a hypothetical newer app
    with pytest.raises(SchemaTooNewError):
        Database(p)

def test_backup_before_migrate_on_upgrade(tmp_path):
    p = str(tmp_path / "t.db")
    db = Database(p)
    db.add_video("vid1", "/a.mp4", "processed", []) # data that must survive
    _set_user_version(p, 3) # simulate an older on-disk schema

    Database(p) # reopen with current code ? upgrades in place, backing up first
    backup = tmp_path / "t.db.bak-v3"
    assert backup.exists(), "an upgrade must snapshot the DB to <db>.bak-v<old> first"
    assert _user_version(p) == SCHEMA_VERSION # upgraded
    assert _user_version(str(backup)) == 3 # the backup is the PRE-upgrade snapshot
    assert Database(p).get_video("vid1") is not None # data preserved across the upgrade

def test_backup_database_export(tmp_path):
    p = str(tmp_path / "t.db")
    db = Database(p)
    db.add_video("vid1", "/a.mp4", "processed", [])
    dest = str(tmp_path / "export" / "my_data.db")
    assert db.backup_database(dest) == dest
    assert os.path.exists(dest) # parent dir created
    assert _user_version(dest) == SCHEMA_VERSION
    assert Database(dest).get_video("vid1") is not None # the export is a usable DB

```

617. user (2026-06-21T19:04:18.861Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p2d\tests\test_db_upgrade_safety.py
(file state is current in your context ? no need to Read it back)

618. user (2026-06-21T19:04:36.828Z)

```

..... [100%]
62 passed in 10.49s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

619. assistant (2026-06-21T19:04:44.251Z)

62 pass (incl. all migration-related suites). Now lint/type + full suite.

620. user (2026-06-21T19:07:24.246Z)

```
=== RUFF+MYPY ===
All checks passed!
Success: no issues found in 1 source file
=== FULL SUITE + COV ===
backend\infrastructure\database.py          388      32      92%
TOTAL                                     9972    1478     85%
Required test coverage of 70% reached. Total coverage: 85.18%
1331 passed, 7 skipped, 7 warnings in 151.20s (0:02:31)
  RESULT: PASS - all tests green
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

621. assistant (2026-06-21T19:07:50.727Z)

Full suite green (1331), coverage 85.18%, `database.py` 92%. P2d has real logic (guard + backup), and you asked for the same diligence ? so a focused review before commit.

622. user (2026-06-21T19:09:48.261Z)

I have everything needed for the review.

Adversarial review: PR P2d (DB upgrade-safety)

****1. DOWNGRADE GUARD**** ? `SCHEMA\_VERSION = 7` (database.py:18) matches the ladder's max (`\_migrate\_to\_v7` / `PRAGMA user\_version = 7` at database.py:395-402; no `\_migrate\_to\_v8`). The pin test `test\_schema\_version\_matches\_ladder` (test_db_upgrade_safety.py:28) plus the existing `==7` tests enforce drift. ****No false-positive risk**** on a current/older DB: the guard is `version > SCHEMA\_VERSION` (database.py:109), so a v7 DB (`7 > 7` false) and any older DB pass straight through. ****No leaked handle on raise****: the guard raises **inside** `with self.\_connect()` as conn: (database.py:96), and `\_connect`'s `finally: conn.close()` (database.py:42) runs during exception unwind ? verified the connection is closed before `SchemaTooNewError` escapes `\_\_init\_\_`. Clean.

****2. BACKUP-BEFORE-MIGRATE**** ? gating `0 < version < SCHEMA\_VERSION` (database.py:117) is correct (fresh=0 skips, current=7 skips). ****WAL snapshot completeness verified empirically****: `PRAGMA wal\_checkpoint(FULL)` (database.py:152) synchronously flushes all WAL frames into the main file **before** `shutil.copy2` of the single main file ? I reproduced this with a 24KB untruncated WAL and even a second open connection, and the copied-main-file-only backup contained every row + the correct `user\_version`, with no orphaned **-wal** sidecar. The single-instance lock (P0e) also rules out a real concurrent writer in production. Copying while `conn` is open is safe (FULL completed first). Best-effort WARN-not-block (database.py:155-157, catches `OSError`/`sqlite3.Error`) is the right call for a downloadable app. `from\_version` is an `int` read from `PRAGMA user\_version` (database.py:104) ? not user-controlled, no injection/path concern in the f-string filename (database.py:150). **.bak-vN** is keyed by `from\_version` so each upgrade-from-version writes exactly one file and re-running overwrites idempotently ? bounded, no unbounded accumulation.

****3. backup_database(dest)**** ? checkpoint runs on the fresh `\_connect` and the FULL checkpoint completes **synchronously inside** `conn.execute` (database.py:166) before the `with` closes, so the flush is durably in the main file by the time `shutil.copy2` (database.py:167) runs after close. `os.makedirs(parent, exist\_ok=True)` guarded by `if parent` (database.py:163-164) handles a bare filename safely. `test\_backup\_database\_export` confirms the export is a usable DB at `SCHEMA\_VERSION`. Returns `dest`. Correct.

****4. INTERACTION**** ? `acquire\_lock` (main.py:958) runs before `Database(db\_path)` (main.py:973), so the guard/backup execute under the single-instance lock as intended ? no concurrent-writer hazard during checkpoint/copy. Fresh-DB path (version 0): guard skipped (`0 > 7` false), backup skipped (`0 < 0` false) ? `test\_fresh\_db\_makes\_no\_backup` + `test\_reopen\_current\_db\_no\_backup\_no\_error` confirm zero behavior change for the normal

fresh/current case.

****5. EDGE**** ? Zero-byte file: reads `user\_version = 0` ? treated as fresh, no guard, no backup (verified). Corrupt non-SQLite file: raises `sqlite3.DatabaseError` from the first statement in `\_init\_db` ? but this is *pre-existing* behavior unchanged by P2d (base-table creation already ran before this PR), and there's nothing to guard/back-up anyway. Read-only FS: WAL pragma already degrades via the `except sqlite3.OperationalError: pass` at database.py:55-56; backup failure degrades via the WARN path (database.py:155). Both degrade gracefully.

Findings

No blocker/major/minor issues found.

- `[nit] backend/infrastructure/database.py:155` ? `\_backup\_before\_migrate` catches `(OSError, sqlite3.Error)`; an unexpected non-OSError/non-sqlite exception during checkpoint/copy (e.g. a `shutil` edge or a `MemoryError`) would propagate and brick `\_\_init\_\_`, defeating the "never block the upgrade" intent. Extremely unlikely in practice given copy2/checkpoint failure modes are OSError/sqlite3.Error. Fix (optional): broaden to `except Exception` for true best-effort, matching the repo's swallow-and-log convention used by the job writers. ? Judgment call; the narrow catch is also defensible.

****Verdict:**** Clean ? the downgrade guard, WAL-complete backup (empirically verified), and export are correct and well-tested; ship it. (One optional nit on the backup exception breadth.)
agentId: a95b4528ddf71d428 (use SendMessage with to: 'a95b4528ddf71d428' to continue this agent)
<usage>subagent_tokens: 72256
tool_uses: 13
duration_ms: 104182</usage>

623. assistant (2026-06-21T19:09:59.056Z)

Clean review ? one optional nit worth taking: broaden the backup's exception catch to `except Exception` so a backup failure can *never* block the upgrade (its whole contract is "WARN, never block"). Quick fix.

624. user (2026-06-21T19:10:04.140Z)

The file C:\Users\avidu\Projects\bhi-wt-p2d\backend\infrastructure\database.py has been updated successfully. (file state is current in your context ? no need to Read it back)

625. user (2026-06-21T19:10:24.976Z)

All checks passed!

.....

[100%]

6 passed in 0.37s

origin/master: f01dc85 feat(db): v7 ? record the UI locale per job (PMF analytics) +

count_jobs_by_locale (#306)

warning: in the working copy of 'tests/test_db_upgrade_safety.py', LF will be replaced by CRLF the next time Git touches it

4176ce5 feat(pkg): P2d ? DB upgrade-safety (downgrade guard + backup-before-migrate + export) (default-OFF)

To <https://github.com/Khelsutra/badminton-highlight-indexer.git>

* [new branch] feat/packaging-p2d-db-safety -> feat/packaging-p2d-db-safety

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

626. assistant (2026-06-21T19:10:45.681Z)

P2d committed cleanly (based on current master `f01dc85`, which has #306's v7 ? hence `SCHEMA\_VERSION=7` is correct). Opening the PR.

627. user (2026-06-21T19:10:57.549Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/307>

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

628. assistant (2026-06-21T19:11:04.926Z)

****PR #307 opened: [feat(pkg): P2d ? DB upgrade-safety](https://github.com/Khelsutra/badminton-highlight-indexer/pull/307).**** Watching CI.

629. user (2026-06-21T19:13:51.796Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    18s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622099
Built wheel smoke (install + launch)      pass    1m10s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27914607592/job/82597622101
Dependency license scan (no copleft in runtime deps)  pass    31s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622125
Full test suite (offline, mocked)      pass    2m35s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27914607592/job/8259762212
Golden fixtures cross-OS (macos-latest) pass    14s    https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27914607592/job/82597622126
Golden fixtures cross-OS (ubuntu-latest)      pass    17s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622140
Golden fixtures cross-OS (windows-latest)      pass    56s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622119
Mypy (lenient baseline) pass    32s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622097
Ruff lint      pass    8s    https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622150
Syntax sweep + import smoke (no GPU stack)      pass    18s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27914607592/job/82597622099
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

630. assistant (2026-06-21T19:13:56.199Z)

All 9 lanes green. Merging P2d and cleaning up.

631. user (2026-06-21T19:14:04.907Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
f01dc85..679e642 master    -> origin/master
=== master ===
679e642 feat(pkg): P2d ? DB upgrade-safety (downgrade guard + backup-before-migrate + export)
(default-OFF) (#307)
=== PR #307 ===
MERGED
```

632. assistant (2026-06-21T19:15:35.536Z)

P2d ****merged**** (master `679e642`) ? ****10 PRs this session****. Updating the tracker, then assessing the next item. There's a real wrinkle for the LIGHT-tier install story I should check: the plan claims "LIGHT = no torch," but `import backend.main` pulls torch at module load (`person\_detector`). Let me size that.

633. user (2026-06-21T19:15:39.951Z)

```
=== module-level 'import torch' (top-level, not inside a def/try) across backend ===
backend/pipeline/detectors/person_detector.py:29:import torch
backend/pipeline/detectors/wasb_infer.py:397:    import torch

=== all torch imports (any indent) ? to gauge lazy-vs-eager ===
backend/config/presets.py:90:        import torch
backend/pipeline/detectors/native_wasb_runner.py:56:_CUDA_PROBE_CODE = "import torch;
```

```
print('torch', torch.__version__, 'cuda', torch.cuda.is_available())"
backend/pipeline/detectors/person_detector.py:29:import torch
backend/pipeline/detectors/wasb_infer.py:397:     import torch
backend/pipeline/detectors/wasb_runner.py:75:         f"python -c \"import torch;
print('torch', torch.__version__, "
```

634. assistant (2026-06-21T19:16:03.889Z)

P2d merged. Key finding: `**only ONE eager module-level import torch**`
(`person_detector.py:29`) ? everything else (`wasb_infer`, `presets`) is already lazy/in-
function. That single import is what forces the LIGHT tier to install torch just to `*import*`
`backend.main`. Making it lazy delivers the plan's "LIGHT = no torch" promise in one contained
change ? a high-value capstone. Let me update the tracker, then read `person_detector.py`.

635. user (2026-06-21T19:16:09.058Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been
updated successfully. (file state is current in your context ? no need to Read it back)

636. user (2026-06-21T19:16:15.879Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been
updated successfully. (file state is current in your context ? no need to Read it back)

637. assistant (2026-06-21T19:16:19.121Z)

Now set up the worktree and read `person_detector.py` to defer the torch import cleanly.

638. user (2026-06-21T19:16:22.935Z)

HEAD is now at 679e642 feat(pkg): P2d ? DB upgrade-safety (downgrade guard + backup-before-
migrate + export) (default-OFF) (#307)
=== person_detector.py: imports + how torch is used ===
"""Person-detection rally cue ? torchvision detector + ByteTrack + motion windowing.

Extracted verbatim from `yolo_hybrid` (fusion P1, `docs/MULTI_SIGNAL_FUSION_PLAN.md`):
the `**person cue is a swappable producer**`, not logic baked into one segmenter. Today
`HybridYoloSegmenter` is the only consumer (it delegates Stage-1 here, byte-identically);
next the fusion segmenter consumes the SAME producer over shuttle-seeded windows, and
later the learned classifier reads its persisted features. Keeping detection here (not in
the segmenter) is what makes "add a cue = register a producer" true.

Behaviour is unchanged from the in-lined version ? only the home moved:

- `lazy_load_model` ? load a permissively-licensed torchvision detector
- `make_tracker` ? a fresh supervision ByteTrack per chunk
- `detect_and_track` ? one frame ? confident persons with persistent IDs
- `extract_action_windows_from_chunk` ? a chunk ? high-motion candidate windows
with the motion stats (`peak/mean_velocity`, `active_frame_count`, `track_id_count`)
that feed the candidate `stats` JSON and the learned fusion features.

The model libs (torchvision, supervision) are imported lazily inside the methods, and the
detector weights load only on first detection (`lazy_load_model`) ? so importing the
segmenter registry never loads a detector model, and tests can mock the seams. (torch + cv2
are top-level imports here; CI's import-smoke stubs them, and they're present on the GPU
machine.) Licensing: torchvision detectors are BSD + COCO weights; ByteTrack is MIT
(ADR-005; ultralytics' AGPL YOLO was dropped).

```
"""
import logging
from typing import Any, Callable, Dict, List, Optional, Tuple

import cv2
import torch

from backend.utils.geometry import get_velocity_normalised
```

```
logger = logging.getLogger(__name__)
```

torchvision detection models are trained on COCO where the "person" category is class 1 (index 0 is the implicit background). NB: ultralytics YOLO used class 0 ? this off-by-one is the classic gotcha when swapping detectors. Keep it named.

```
_PERSON_CLASS_ID = 1
```

**Permissively-licensed torchvision detectors selectable via config
`indexing.detector\_model`. All are BSD-licensed (code + COCO weights).**

```
_DETECTOR_FACTORIES = {
    "fasterrcnn_resnet50_fpn": "fasterrcnn_resnet50_fpn",
    "fasterrcnn_mobilenet_v3_large_fpn": "fasterrcnn_mobilenet_v3_large_fpn",
    "retinanet_resnet50_fpn": "retinanet_resnet50_fpn",
    === all 'torch' references in the file ===
1: """Person-detection rally cue ? torchvision detector + ByteTrack + motion windowing.
11: - `lazy_load_model` ? load a permissively-licensed torchvision detector
18: The model libs (torchvision, supervision) are imported lazily inside the methods, and the
20: segmenter registry never loads a detector model, and tests can mock the seams. (torch + cv2
22: machine.) Licensing: torchvision detectors are BSD + COCO weights; ByteTrack is MIT
29: import torch
35: # torchvision detection models are trained on COCO where the "person" category is
40: # Permissively-licensed torchvision detectors selectable via config
53:     Owns the torchvision model + per-chunk ByteTrack tracker. Constructed with the run
60:         # Dynamic torchvision module (None until lazy_load_model); typed Any so the
69:         # Imported lazily so the module imports without torchvision present and so
71:         from torchvision.models import detection as tv_detection
75:         self.device = "cuda" if use_gpu and torch.cuda.is_available() else "cpu"
76:         if use_gpu and not torch.cuda.is_available():
85:         logger.info(f"Loading torchvision detector '{model_name}' on {self.device}...")
106:         detection AND association: torchvision does the detection, supervision's
111:         # cv2 frames are BGR uint8 HxWxC; torchvision wants RGB float [0,1] CxHxW.
113:         tensor = torch.from_numpy(rgb).permute(2, 0, 1).float().div(255.0).to(self.device)
114:         with torch.no_grad():
```

639. assistant (2026-06-21T19:16:37.263Z)

`torch` is used only inside two methods (`lazy\_load\_model`, `detect\_and\_track`) ? torchvision is *already* lazy. So deferring the single top-level `import torch` is clean. cv2 stays eager (opencv-python is a hard dep, always installed). Let me read the two methods to place the lazy imports.

640. user (2026-06-21T19:16:38.537Z)

```
66         if self.model is not None:
67             return
68
69         # Imported lazily so the module imports without torchvision present and so
70         # the test suite (which pre-sets self.model) never pulls in model weights.
71         from torchvision.models import detection as tv_detection
72
73         idx_cfg = self.config.get("indexing", {})
74         use_gpu = idx_cfg.get("use_gpu", True)
75         self.device = "cuda" if use_gpu and torch.cuda.is_available() else "cpu"
76         if use_gpu and not torch.cuda.is_available():
77             logger.warning("use_gpu is True but CUDA is not available. Falling back to
CPU.")
78
79         model_name = idx_cfg.get("detector_model", _DEFAULT_DETECTOR)
80         if model_name not in _DETECTOR_FACTORIES:
81             logger.warning(f"Unknown detector_model '{model_name}'. Falling back to
{_DEFAULT_DETECTOR}.")
```



```

82         model_name = _DEFAULT_DETECTOR
83
84         builder = getattr(tv_detection, _DETECTOR_FACTORIES[model_name])
85         logger.info(f"Loading torchvision detector '{model_name}' on {self.device}...")

```

641. user (2026-06-21T19:16:43.694Z)

```

104         Returns a list of (track_id, bbox) where bbox is (x1, y1, x2, y2) in pixels.
105         This replaces the old ultralytics ``model.track(...)`` call, which bundled
106         detection AND association: torchvision does the detection, supervision's
107         ByteTrack does the association.
108         """
109         import supervision as sv
110
111         # cv2 frames are BGR uint8 HxWxC; torchvision wants RGB float [0,1] CxHxW.
112         rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
113         tensor = torch.from_numpy(rgb).permute(2, 0,
114         1).float().div(255.0).to(self.device)
115         with torch.no_grad():
116             outputs = self.model([tensor])

```

642. assistant (2026-06-21T19:16:51.473Z)

Both methods already lazy-import their model libs (torchvision/supervision) ? I'll add `import torch` alongside. Removing the eager top-level import + updating the docstring.

643. user (2026-06-21T19:16:53.511Z)

The file C:\Users\avidu\Projects\bhi-wt-p2e\backend\pipeline\detectors\person_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

644. user (2026-06-21T19:16:59.451Z)

The file C:\Users\avidu\Projects\bhi-wt-p2e\backend\pipeline\detectors\person_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

645. user (2026-06-21T19:17:05.047Z)

The file C:\Users\avidu\Projects\bhi-wt-p2e\backend\pipeline\detectors\person_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

646. user (2026-06-21T19:17:13.969Z)

The file C:\Users\avidu\Projects\bhi-wt-p2e\backend\pipeline\detectors\person_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

647. assistant (2026-06-21T19:17:38.523Z)

Now the proof test ? a subprocess with `torch`/`torchvision` blocked from import, asserting `backend.main` + `person\_detector` still import (the LIGHT-tier-no-torch guarantee).

648. assistant (2026-06-21T19:17:47.789Z)

```

"""The LIGHT tier imports WITHOUT torch installed (PACKAGING_PLAN.md P2e).

```

```

`backend.main` and the detector registry must import even when torch/torchvision are absent ?
the
LIGHT tier (Gemini / motion / cpu-mock) ships no torch. The person-detection cue is HEAVY-only
and
never instantiated under LIGHT, so its torch usage is deferred into its methods. This spawns a
fresh
interpreter with torch/torchvision BLOCKED at the import machinery to prove it.
"""

```

```

import os
import subprocess
import sys

REPO_ROOT = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))

_SCRIPT = r"""
import sys, importlib.abc

class _BlockTorch(importlib.abc.MetaPathFinder):
    def find_spec(self, name, path, target=None):
        top = name.split('.')[0]
        if top in ('torch', 'torchvision'):
            raise ModuleNotFoundError(f"blocked for the no-torch test: {name}")
        return None

sys.meta_path.insert(0, _BlockTorch())

```

Hard assert torch truly can't be imported in this interpreter.

```

try:
    import torch # noqa: F401
    print("FAIL: torch was importable despite the block"); sys.exit(2)
except ModuleNotFoundError:
    pass

```

The LIGHT-tier import surface must hold without torch:

```

import backend.main # noqa: F401 (must not import torch at module load)
import backend.pipeline.detectors.person_detector # noqa: F401 (torch deferred into methods)
from backend.pipeline.segmenters.base import segmenter_registry
assert "motion" in segmenter_registry.list_available()
print("IMPORT_OK_NO_TORCH")
"""

def test_light_tier_imports_without_torch():
    env = {**os.environ, "PYTHONPATH": REPO_ROOT + os.pathsep + os.environ.get("PYTHONPATH",
    "")}
    proc = subprocess.run(
        [sys.executable, "-c", _SCRIPT],
        cwd=REPO_ROOT, env=env, capture_output=True, text=True, timeout=120,
    )
    assert proc.returncode == 0 and "IMPORT_OK_NO_TORCH" in proc.stdout, (
        "backend.main / person_detector must import with torch absent (LIGHT tier).\n"
        f"stdout:\n{proc.stdout[-2000:]} \nstderr:\n{proc.stderr[-2000:]}")

```

649. user (2026-06-21T19:17:47.818Z)

File created successfully at: C:\Users\avidu\Projects\bhi-wt-p2e\tests\test_light_tier_no_torch.py (file state is current in your context ? no need to Read it back)

650. user (2026-06-21T19:17:58.426Z)

```

. [100%]
1 passed in 4.54s
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

```

651. assistant (2026-06-21T19:18:05.261Z)

LIGHT tier now genuinely torch-free (proof passes). Now confirm the HEAVY/yolo path (which *uses* person_detector) still works ? full suite + lint.

652. user (2026-06-21T19:20:47.945Z)

=== person-detector / yolo tests ===

```
..\..\AppData\Local\Programs\Python\Python313\Lib\unittest\mock.py:1473: AttributeError
===== short test summary info =====
FAILED
tests/test_person_detector.py::test_detect_and_track_filters_persons_and_returns_track_ids
1 failed, 91 passed, 1247 deselected in 16.30s
=== ruff+mypy ===
All checks passed!
Success: no issues found in 1 source file
=== FULL SUITE ===
1 failed, 1331 passed, 7 skipped in 94.36s (0:01:34)
=====
RESULT: FAIL - pytest exit 1
=====
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

653. assistant (2026-06-21T19:20:56.244Z)

One test failed ? it patches `person\_detector.torch` (the module attribute I removed). I need to see how it mocks torch and adapt it to the now-lazy import.

654. user (2026-06-21T19:20:57.660Z)

```
3:Pure/offline ? the torchvision model and supervision tracker are mocked, so these run
4:without GPU/torch weights. They pin the COCO person==1 filter (the YOLO-was-0 gotcha) and
7:from unittest.mock import MagicMock, patch
12:def _make_cap_mock(num_frames=300):
13:    """A cv2.VideoCapture mock yielding `num_frames` readable frames at 30fps."""
15:    cap_mock = MagicMock()
16:    cap_mock.isOpened.return_value = True
18:    def mock_get_prop(prop):
25:    cap_mock.get.side_effect = mock_get_prop
26:    cap_mock.read.side_effect = [(True, "frame")] * num_frames + [(False, None)]
27:    return cap_mock
36:@patch("backend.pipeline.detectors.person_detector.cv2.VideoCapture")
37:def test_extract_action_windows_returns_enriched_dicts(mock_cap):
42:    mock_cap.return_value = _make_cap_mock()
60:@patch("backend.pipeline.detectors.person_detector.cv2.VideoCapture")
61:def test_detections_per_frame(mock_cap):
71:    cap = _make_cap_mock(num_frames=20)
83:    mock_cap.return_value = cap
93:@patch("backend.pipeline.detectors.person_detector.cv2.VideoCapture")
94:def test_detections_over_windows_single_pass(mock_cap):
95:    """One sequential pass dispatches sampled frames to the right windows (no per-window
seek)."""
109:    mock_cap.return_value = cap
118:    mock_cap.assert_called_once() # opened ONCE for all windows
121:def test_detect_and_track_filters_persons_and_returns_track_ids():
123:    tracker, and returns (track_id, bbox) tuples ? exercising the real torchvision
130:    # Fake torchvision detector output: a person (kept), a low-score person (dropped),
149:    with patch("backend.pipeline.detectors.person_detector.cv2.cvtColor",
return_value=np.zeros((10, 10, 3))), \
150:        patch("backend.pipeline.detectors.person_detector.torch") as mock_torch, \
151:        patch.dict("sys.modules", {"supervision": MagicMock()}):
152:        mock_torch.no_grad.return_value.__enter__ = lambda s: None
153:        mock_torch.no_grad.return_value.__exit__ = lambda s, *a: None
154:        # torch.from_numpy(...).permute(...).float().div(...).to(...) -> dummy tensor
155:        mock_torch.from_numpy.return_value.permute.return_value.float.return_value.div.retur
n_value.to.return_value = "tensor"
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

655. assistant (2026-06-21T19:21:07.993Z)

The test patches `person\_detector.torch` (the removed module attribute); since `torch` is now imported inside the method, it should be injected via `sys.modules` instead. Let me read the full test and fix it.

656. user (2026-06-21T19:21:08.962Z)

```

121     def test_detect_and_track_filters_persons_and_returns_track_ids():
122         """detect_and_track keeps only confident PERSON detections, feeds them to the
123         tracker, and returns (track_id, bbox) tuples ? exercising the real torchvision
124         output-shape handling and the COCO person==1 filter (the YOLO-was-0 gotcha)."""
125         import numpy as np
126
127         pd = PersonDetector({"indexing": {}})
128         pd.device = "cpu"
129
130         # Fake torchvision detector output: a person (kept), a low-score person (dropped),
131         # and a racket/other class (dropped).
132         out = {"boxes": MagicMock(), "labels": MagicMock(), "scores": MagicMock()}
133         out["boxes"].detach.return_value.cpu.return_value.numpy.return_value = np.array(
134             [[0, 0, 10, 10], [5, 5, 15, 15], [20, 20, 30, 30]], dtype=float)
135         out["labels"].detach.return_value.cpu.return_value.numpy.return_value = np.array(
136             [_PERSON_CLASS_ID, _PERSON_CLASS_ID, 38])          # 38 = sports-ball-ish, not
person
137         out["scores"].detach.return_value.cpu.return_value.numpy.return_value = np.array(
138             [0.9, 0.1, 0.99])
139         pd.model = MagicMock(return_value=out)
140
141         # Fake tracker echoes the detections it receives, assigning a single track id.
142         tracker = MagicMock()
143         tracked = MagicMock()
144         tracked.__len__.return_value = 1
145         tracked.tracker_id = [7]
146         tracked.xyxy = np.array([[0, 0, 10, 10]], dtype=float)
147         tracker.update_with_detections.return_value = tracked
148
149         with patch("backend.pipeline.detectors.person_detector.cv2.cvtColor",
return_value=np.zeros((10, 10, 3))), \
150             patch("backend.pipeline.detectors.person_detector.torch") as mock_torch, \
151             patch.dict("sys.modules", {"supervision": MagicMock()}):
152             mock_torch.no_grad.return_value.__enter__ = lambda s: None
153             mock_torch.no_grad.return_value.__exit__ = lambda s, *a: None
154             # torch.from_numpy(...).permute(...).float().div(...).to(...) -> dummy tensor
155             mock_torch.from_numpy.return_value.permute.return_value.float.return_value.div.r
eturn_value.to.return_value = "tensor"
156             result = pd.detect_and_track("frame", tracker, score_thresh=0.5)
157
158             # Only the high-confidence person survived the score+class filter
159             assert result == [(7, (0.0, 0.0, 10.0, 10.0))]
160             # And exactly that filtered set (1 box) was handed to the tracker
161             tracker.update_with_detections.assert_called_once()
162

```

657. user (2026-06-21T19:21:18.773Z)

The file C:\Users\avidu\Projects\bhi-wt-p2e\tests\test_person_detector.py has been updated successfully. (file state is current in your context ? no need to Read it back)

658. user (2026-06-21T19:22:56.592Z)

```

.....
5 passed in 3.35s
=== FULL SUITE ===

```

[100%]

1332 passed, 7 skipped in 85.76s (0:01:25)

=====

RESULT: PASS - all tests green

=====

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

659. assistant (2026-06-21T19:23:11.570Z)

Full suite green (1332). The HEAVY/yolo path still works (test adapted to the lazy import), and the LIGHT tier imports torch-free. Committing P2e.

660. user (2026-06-21T19:23:25.925Z)

All checks passed!

origin/master: 679e642 feat(pkg): P2d ? DB upgrade-safety (downgrade guard + backup-before-migrate + export) (default-OFF) (#307)

warning: in the working copy of 'tests/test_light_tier_no_torch.py', LF will be replaced by CRLF the next time Git touches it

252926a feat(pkg): P2e ? lazy-torch so the LIGHT tier imports with NO torch installed (default-OFF)

To <https://github.com/Khelsutra/badminton-highlight-indexer.git>

* [new branch] feat/packaging-p2e-lazy-torch -> feat/packaging-p2e-lazy-torch

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

661. user (2026-06-21T19:23:45.259Z)

<https://github.com/Khelsutra/badminton-highlight-indexer/pull/308>

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

662. assistant (2026-06-21T19:23:49.419Z)

PR #308 opened: [feat(pkg): P2e ? lazy-torch (torch-free LIGHT tier)](<https://github.com/Khelsutra/badminton-highlight-indexer/pull/308>). Watching CI.

663. user (2026-06-21T19:26:36.136Z)

Syntax sweep + import smoke (no GPU stack) pass 20s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598456017>

Built wheel smoke (install + launch) pass 1m11s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598455958>

Dependency license scan (no copyleft in runtime deps) pass 30s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598455997>

Full test suite (offline, mocked) pass 2m31s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598455963>

Golden fixtures cross-OS (macos-latest) pass 17s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598456025>

Golden fixtures cross-OS (ubuntu-latest) pass 22s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598455971>

Golden fixtures cross-OS (windows-latest) pass 57s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598455977>

Mypy (lenient baseline) pass 31s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598456060>

Ruff lint pass 9s <https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598455985>

Syntax sweep + import smoke (no GPU stack) pass 20s

<https://github.com/Khelsutra/badminton-highlight-indexer/actions/runs/27914920652/job/82598456017>

Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

664. assistant (2026-06-21T19:26:40.514Z)

All 9 lanes green. Merging P2e and cleaning up.

665. user (2026-06-21T19:26:54.224Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
679e642..de2bf14 master -> origin/master
=== master ===
de2bf14 feat(pkg): P2e ? lazy-torch so the LIGHT tier imports with NO torch installed (default-OFF) (#308)
=== PR #308 ===
MERGED
=== open PRs (mine) ===
=== no VMs (cost audit) ===
Listed 0 items.
```

666. assistant (2026-06-21T19:27:26.582Z)

P2e ****merged**** (master `de2bf14`) ? ****11 PRs this session****, 0 open PRs, 0 VMs (\$0). The LIGHT tier is now genuinely torch-free. Updating the tracker + memory, then I'll checkpoint with a grounded recommendation.

667. user (2026-06-21T19:27:30.346Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

668. user (2026-06-21T19:27:46.479Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\packaging-app-plan.md has been updated successfully. (file state is current in your context ? no need to Read it back)

669. assistant (2026-06-21T19:28:15.752Z)

This is a clean, comprehensive checkpoint. Here's where everything stands.

Session tally: 11 PRs merged + the plan

`origin/master` = `de2bf14`, ****0 open PRs****, ****0 VMs (\$0 spend)****, all worktrees cleaned.

```
**P0 ? engine prerequisites (5):** app-home `#293` · `/api/config` redaction `#295` · bind + DNS-rebinding guard `#296` · scipy + license-scan & built-wheel CI lanes `#297` · single-instance lock `#298`
**P1 ? UI in the box (2):** console/wizard endpoints `#300` · **SPA bundled single-origin into the wheel** `#302`
**P2 ? the app (4):** **`indexer --app` launcher** `#303` · ffmpeg resolver `#305` · DB upgrade-safety `#307` · **lazy-torch (LIGHT tier now torch-free)** `#308`
```

Every PR: isolated worktree off `origin/master` (which moved ~6 more times under me ? rebased cleanly through #304/#306), explicit-path commits, adversarial review on the logic PRs, full suite + 9-lane CI green, coverage held ~85%.

****Net result:**** a user can `pip install` the engine, get the real KhelSutra UI ****in the wheel****, run ****`indexer --app`**** ? it boots and opens their browser to the bundled SPA ? ****with no torch**, no manual ffmpeg setup, safe app-home state, and safe DB upgrades.****** The LIGHT-tier "batteries-included" app is engine-complete.

On "through P3 with GPU + screencasts" ? what I did and why

You chose ****option 1**** for the GPU run (execute it fresh), and I continued the engine items (P2a/P2d/P2e). I'm now ****11 PRs deep****, and the two remaining headline items ? the ****GPU/P3 validation**** and the ****visual screencast**** ? are exactly the kind of work I shouldn't start at the tail of a long session:

- **The GPU VM is billable** and must be torn down start-to-finish; abandoning one mid-op because context ran out is the precise cost mistake `[[gcp-vm-always-delete]]` guards against. It's **verified ready** (authed, \$0 now, scripts + corpus staged, ~\$0.50/run) with a **known native-WASB 0-rally config blocker** to chase ? it deserves a focused push.

- **The screencast** needs a live browser-capture session; it's most valuable produced once, ideally *on* the GPU box (real end-to-end on the target environment).

I've recorded the **exact GPU plan** (provision cmd · L4 cost · validate-the-wheel + native-WASB · screencast · DELETE + confirm \$0) in memory so a fresh push starts it immediately and safely.

Recommended next (fresh, focused)

1. **P3 GPU validation + screencast** ? provision an L4, install the *wheel* on the real GPU box, validate the full packaged stack (``indexer --app`` ? bundled SPA ? native-WASB ? reel), screencast it, delete the VM. (Also resolve the 22-rally tuned-config reproducibility.)
2. **P2b-install** ? ``uv`/generated cross-platform install script` (now a clean **torch-free** LIGHT install thanks to P2e) + a truly-local default ``config.json`` + uninstall manifest.
3. **P2c wizard** ? first-run UI in the khelsutra SPA (cross-repo, coordinated like P1b).

I've stopped here rather than half-do the GPU run ? say the word in a fresh session and I'll execute the GPU validation end-to-end with teardown as priority #1, or pick up P2b-install if you'd rather keep the no-spend engine work going.

670. user (2026-06-21T21:46:06.140Z)

can you capture all the Recommended next steps in an icebreaker prompt, ensure teh docs and handoffs are all fresh so that I can start this work in a new session? Amazing work and got a lot done by you!

671. assistant (2026-06-21T21:47:17.756Z)

Thank you ? that was a genuinely big push. Let me make everything fresh and findable for a clean new-session start: finalize the tracked doc, commit it to the repo, refresh the memory handoffs, and write you a paste-ready icebreaker prompt. First, let me read the plan doc's current header/tracker to finalize it.

672. user (2026-06-21T21:47:20.598Z)

Batteries-Included App ? Packaging & Distribution Plan

```
> Status: `DRAFT` ? 01/02 resolved 2026-06-21 (LIGHT-first; cross-platform webserver, no native component; WASB-complete north star); 03+ remain owner/counsel-gated · Owner: Avi Dullu · Created: 2026-06-21 · Last updated: 2026-06-21
> Lifecycle: `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/` when DONE)
> Tracking anchors: $7 progress tracker is the source of truth; index in `docs/README.md`; pointer in memory `current-state` + `docs/NEXT_STEPS.md` once work starts.
> Relation to existing docs: supersedes the narrative of `docs/DESKTOP_APP_PLAN.md` (per locked D10 ? self-hosted webserver, not a desktop app); builds on `docs/COMPUTE_DECOUPLED_SERVING/` (the profile/storage/compute seam) and is the distribution layer under `khelsutra/docs/LOCAL_APPLIANCE_ARCHITECTURE.md`.
> Honesty note: claims tagged `[verified]` (checked at file:line this session), `[researched]` (web/needs sign-off), `[design]` (proposed). Not legal advice ? licensing items need counsel.
```

0. TL;DR

Ship the engine as **ONE pip-installable artifact** whose behaviour is selected entirely by the already-merged decoupled-compute seam (``deployment.profile`` / ``compute_target`` / Storage+Compute backends), fronted by a friendly launcher + a Windows installer ? **not** a frozen monolith, **not** an Electron/Tauri shell. The hard constraint that forces scope: torch is deliberately out of ``pyproject`` (needs ``--index-url``) **and** native WASB needs a *second, incompatible* torch env (conda py3.8 + torch 1.11+cu113) shelled out via ``$WASB_PYTHON`` ``[verified native_wasb_runner.py:89]`` ? so no single venv/freeze can deliver the GPU-quality

path. ****The honest v1 is a LIGHT tier**** (Gemini + motion + cpu-mock, ***zero*** torch/GPU) that preserves CUJ-1.5/8; the native-WASB GPU path (CUJ-6/7) ships ****owner-only**** until two blockers clear (WASB-C3 weight provenance + a reproducible truly-local config). ****LIGHT is the FIRST shippable increment, not the terminal scope ? the sole completion goal (north star, L10) is the FULL quality path: native WASB + BYO training + in-app annotation.**** Delivery is ****cross-platform (Windows/macOS/Linux): the same local webserver powers the UI on all three, with NO native component**** (L9). ****Five small engine prerequisites must land before any packaging push**** (\$7 P0). The single most actionable finding: ****all mutable state is CWD-relative today and `--setup`/server disagree on `config.json`**** ? a silent data-loss footgun that P0 must fix first.

...

5:> ****Tracking anchors:**** \$7 progress tracker is the source of truth; index in `docs/README.md`; pointer in memory `current-state` + `docs/NEXT\_STEPS.md` once work starts.

89:| ****P0a**** | ****App-data HOME**** (`RALLY\_APP\_HOME` env + per-OS default via `default\_os\_app\_home`) threaded through `default\_config\_path()`,

`load\_dotenv(env\_file\_to\_load())`, `resolve\_output\_dir` (output/DB), package-relative static mount; ****unified `doctor.py` write with the server read****. Byte-identical when unset; 4-lens adversarial review folded. | ? | No | ? | [#293](https://github.com/Khelsutra/badminton-highlight-indexer/pull/293) |

90:| P0b | ****Redact `GET /api/config`**** (`backend/utls/redact.py`) + serialization-safety test (incl. a full-`AppConfig` over-masking guard). ***(Host-header guard moved to P0c ? same exposure concern.)*** | ? | No | ? | [#295](https://github.com/Khelsutra/badminton-highlight-indexer/pull/295) |

91:| P0c | ****Config/profile-driven bind + DNS-rebinding Host-header guard**** (`TrustedHostMiddleware` + `\_resolve\_bind\_host`; default `127.0.0.1`; `0.0.0.0`+`RALLY\_ALLOWED\_HOSTS` under edge-auth) + a startup WARN for the edge-auth/allowlist footgun. Byte-identical default; adversarial review folded. | ? | No | ? | [#296](https://github.com/Khelsutra/badminton-highlight-indexer/pull/296) |

92:| P0d | ****Declare `scipy`**** in `[project.dependencies]` (serving use `[verified features/rally\_seq.py:17]`, was transitive-only) + ****license-scan CI lane**** (pip-licenses; fails on GPL/AGPL/LGPL/SSPL) + ****built-wheel smoke lane**** (build ? install ? `indexer --help` from a non-repo cwd). Both lanes green on CI. | ? | No | ? |

[#297](https://github.com/Khelsutra/badminton-highlight-indexer/pull/297) |

93:| P0e | ****Single-instance OS lock**** (`backend/utls/single\_instance.py`, fcntl/msvcrt) on `.lock` before `fail\_stale\_jobs` ? a 2nd instance on the same DB refuses; unwritable-path degrades to WARN. Subprocess two-server proof. | P0a | No | ? |

[#298](https://github.com/Khelsutra/badminton-highlight-indexer/pull/298) |

95:*** P0 (engine-prerequisite tranche) COMPLETE ? P0a?P0e all merged 2026-06-21.** Each isolated, default-OFF, adversarially reviewed, full-suite + 9-lane CI green; coverage held ~85%. The engine is now safe to launch from an arbitrary directory and to bundle. ****? NEXT = P1**** (bundle the SPA single-origin + the 3 net-new endpoints `/api/capabilities`, `/api/videos` list, `/api/reels` + `/api/version` contract test).

96:| ****P1a**** | ****Net-new endpoints**** (engine-only): `GET /api/version` + `/api/capabilities` (segmenters/ffmpeg/nvidia-smi/gemini-present/weights/profile ? presence-bools only) + `/api/videos` (list; new `Database.list\_videos`) + `/api/reels` (+ `/api/reels/{filename}` download; ****global flat list**** ? reels have no per-video assoc today, honest scope). | ? | No | ? | [#300](https://github.com/Khelsutra/badminton-highlight-indexer/pull/300) |

97:| P1b | ****Bundle the SPA single-origin****: source-agnostic `scripts/bundle\_ui.sh` (build `khelsutra/web`@pinned now; unpack a release tarball later ? option 3) ? committed `backend/ui` snapshot (option 1, ~700KB) + `ui\_version.json` + `THIRD\_PARTY\_NOTICES.md`; hatchling auto- includes it (no force-include); FastAPI serves `/` + `/assets` (legacy fallback); `/api/version` reports the bundled UI; SPA?engine skew-guard contract test; `backend/ui` **** -text`**. | P1a | No (cross-repo) | ? | [#302](https://github.com/Khelsutra/badminton-highlight-indexer/pull/302) |

99:*** P1 COMPLETE (P1a #300 + P1b #302) ? the engine now serves the real KhelSutra SPA single-origin from the wheel, with the console/wizard endpoints in place. ? NEXT = P2** (cross-platform launcher `indexer --app` + first-run wizard rendered from `/api/capabilities`). P2's wizard UI is SPA work (khelsutra repo) ? coordinate like P1b.

100:| ****P2a**** | ****ffmpeg/ffprobe resolver**** (`backend/utls/ffmpeg.py`: `ffmpeg\_bin`/`ffprobe\_bin` env-overridable via `RALLY\_FFMPEG`/`RALLY\_FFPROBE` + `require\_ffmpeg` fail-loud w/ per-OS hint) wrapping all 13 argv sites; byte-identical when unset. ***(first-run cpu-mock self-test deferred to the wizard slice.)*** | ? | No | ? |

[#305](https://github.com/Khelsutra/badminton-highlight-indexer/pull/305) |

101:| P2b-launcher | ****Pure-Python `indexer --app` launcher**** ? (Win/Mac/Linux, no native component): localhost-only bind + free-port fallback + health-poll ? `webbrowser.open` (`RALLY\_NO\_BROWSER` escape); never CLI single-shot; rides the P0e lock + P0a app-home. | P1a |


```
No | ? | [#303](https://github.com/Khelsutra/badminton-highlight-indexer/pull/303) |
102:| P2b-install | **Install path** (`uv tool install` / a generated cross-platform install
script; ship `config.json` defaulting truly-local; optional thin per-OS shortcut to `indexer
--app`) | P2b-launcher | No (signing only if clickable binaries ? 09) | ? | ? |
103:| P2c | **In-UI 3-step first-run wizard** (Gemini key / Compute / Videos), rendered from
`/api/capabilities`, en/kn/hi | P1b | No | ? | ? |
104:| P2d | **DB upgrade contract** ? `SCHEMA_VERSION` + downgrade guard (`SchemaTooNewError`) +
backup-before-migrate (`<db>.bak-v<old>`, WAL-complete) + `backup_database()` export. Byte-
identical for fresh/current DBs. *(uninstall manifest deferred to P2b-install.)* | ? | No | ? |
[#307](https://github.com/Khelsutra/badminton-highlight-indexer/pull/307) |
105:| P2e | **Lazy-torch for the LIGHT tier** ? deferred the one eager `import torch`
(`person_detector.py`) so `import backend.main` works WITHOUT torch (the LIGHT-tier promise,
enforced by a blocked-torch subprocess test). | ? | No | ? |
[#308](https://github.com/Khelsutra/badminton-highlight-indexer/pull/308) |
106:| **P3a** | **micromamba + conda-forge** replaces Miniconda+`conda tos accept` in
`deploy/cloudrun/Dockerfile` (+ ADR); reuse for the local managed WASB env | ? | No | ? | ? |
```

673. assistant (2026-06-21T21:47:31.289Z)

The doc is current and accurate. Let me finalize it (status ? IN PROGRESS, a "? NEXT SESSION" block), then commit it to the repo so a new session finds it. First the status header + a next-session block.

674. user (2026-06-21T21:47:54.015Z)

The file C:\Users\avidu\Projects\badminton-highlight-indexer\docs\PACKAGING_PLAN.md has been updated successfully. (file state is current in your context ? no need to Read it back)

675. assistant (2026-06-21T21:48:02.415Z)

Now commit the finalized doc to the repo (so a new session finds it) + add index pointers. Worktree + copy + read the index files.

676. user (2026-06-21T21:48:06.589Z)

```
HEAD is now at de2bf14 feat(pkg): P2e ? lazy-torch so the LIGHT tier imports with NO torch
installed (default-OFF) (#308)
=== copied; size ===
169 C:/Users/avidu/Projects/bhi-wt-docs/docs/PACKAGING_PLAN.md
=== docs/README.md index area (where to add the pointer) ===
13:## ? Documentation status & governance
16:## Product / platform plans
19:- [DESKTOP_APP_PLAN.md](DESKTOP_APP_PLAN.md) ? **Tracked scoping doc (2026-06-18, owner-
gated, `[design]` ? NOTHING built):** package the local pipeline as a **cross-platform desktop
app** for non-technical players ? plug in the camera's USB/SD card, click once, get back **only
the highlight reels**, fully on-device (no upload, no original-video bloat). The centerpiece is
the **de-WSL native-compute port** (run WASB natively per-OS: Linux/Windows CUDA · macOS MPS/CPU
? drop WSL2), and its **make-or-break unknown**: is CPU acceptable on a typical *no-GPU*
enthusiast laptop (WASB is ~3.4× real-time on a laptop 2070S)? $7 tracker = **P0 compute-
feasibility spike ? P1 native detector ? P2 app shell ? P3 packaging/installers**. Realizes the
**L0 fully-local** tier of [VIDEO_LOCALITY_MODEL.md](VIDEO_LOCALITY_MODEL.md); is the
**`LocalDesktop` ComputeBackend** of [PLATFORM_ARCHITECTURE.md](PLATFORM_ARCHITECTURE.md) $3.2;
the smaller/export-clean owned model
([OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md)) is one answer to the
compute crux.
21:## Practical / ongoing
26:## Data layer (collection · annotation · golden sets ? **no ML**)
29:## Conventions & templates
32:## Historical documents
```

677. assistant (2026-06-21T21:48:11.823Z)

Let me read the exact README insertion point + the NEXT_STEPS top.

678. user (2026-06-21T21:48:12.792Z)

```
16     ## Product / platform plans
17     - [COMPUTE_DECOUPLED_SERVING/](COMPUTE_DECOUPLED_SERVING/README.md) ? *** Active tracked
project (2026-06-21, `DRAFT` owner-gated):** the productionization successor to
DECOUPLED_COMPUTE (its deferred M3/M4). **One pure core (DETECT?WINDOW?STITCH) that never
branches on deployment profile**, with a `ComputeBackend` (where it runs) + `StorageBackend`
(what bytes) selected by **one startup config**. Ships **truly-local** (local/USB + local GPU,
in-process stitch, zero cloud) + **cloud-serving** (upload<2GB / gdrive / bucket ? a **Cloud Run
GPU job runs detect AND stitch**, metered into a per-job cost ledger) + **cpu-mock** (CI).
Includes [TESTING_STRATEGY.md](COMPUTE_DECOUPLED_SERVING/TESTING_STRATEGY.md) (profile-PARITY
proof for $0), [architecture.svg](COMPUTE_DECOUPLED_SERVING/architecture.svg), a
[runlogs/](COMPUTE_DECOUPLED_SERVING/runlogs/) posterity dir, and **ADR-014** ([lineage
ADR-009?010?011?012?013?014](archives/decisions/DECISIONS.md)). $7 M0?M10 tracker is the source
of truth.
18     - [DECOUPLED_COMPUTE/](archives/past_projects/DECOUPLED_COMPUTE/README.md) ? ?
**DELIVERED + ARCHIVED (2026-06-21).** **Lifted the rally pipeline onto a **rented cloud-native
GPU** doing bucket-driven `train` + `infer`: storage/compute seams (PRs #246?#259) + the **M3.5
native WASB runner**; **M0?M2 + M3.5 proven on a GCP L4** (M2 ? `weights/0.1.0-l4/`). Delivered
on **GCP** ? DigitalOcean dropped for compute ([ADR-013](archives/decisions/DECISIONS.md)).
**M3/M4 (serving endpoint + per-video billing + scale-to-zero) deferred ? the "Cloud Run + GPU
serving" subproject** ([ADR-012](archives/decisions/DECISIONS.md)). Snapshot under
`docs/archives/past_projects/DECOUPLED_COMPUTE/` (incl. `DEPLOY_REPORT_GCP_2026-06-20.md`).
19     - [DESKTOP_APP_PLAN.md](DESKTOP_APP_PLAN.md) ? **Tracked scoping doc (2026-06-18, owner-
gated, `[design]` ? NOTHING built):** package the local pipeline as a **cross-platform desktop
app** for non-technical players ? plug in the camera's USB/SD card, click once, get back **only
the highlight reels**, fully on-device (no upload, no original-video bloat). The centerpiece is
the **de-WSL native-compute port** (run WASB natively per-OS: Linux/Windows CUDA · macOS MPS/CPU
? drop WSL2), and its **make-or-break unknown**: is CPU acceptable on a typical *no-GPU*
enthusiast laptop (WASB is ~3.4x real-time on a laptop 2070S)? $7 tracker = **P0 compute-
feasibility spike ? P1 native detector ? P2 app shell ? P3 packaging/installers**. Realizes the
**L0 fully-local** tier of [VIDEO_LOCALITY_MODEL.md](VIDEO_LOCALITY_MODEL.md); is the
**`LocalDesktop` ComputeBackend** of [PLATFORM_ARCHITECTURE.md](PLATFORM_ARCHITECTURE.md) §3.2;
the smaller/export-clean owned model
([OWNED_MODEL_IMPLEMENTATION_PLAN.md](OWNED_MODEL_IMPLEMENTATION_PLAN.md)) is one answer to the
compute crux.
20
```

679. user (2026-06-21T21:48:13.622Z)

```
1     # Next steps (live) ? refreshed 2026-06-21
2
3     > **How to read:** this top section is the **current cross-track state + prioritized
pending list** (refreshed 2026-06-21, replacing the stale 2026-06-13 header). The detailed
**owned-model & rally-quality track** plan (2026-06-13) is **preserved below, still valid** ? it
is now *one of several* active tracks. Live blow-by-blow: memory `current-state.md`. Doc
lifecycle/hygiene: `docs/DOC_STATUS.md`.
4
5     ## ? Where we are (2026-06-21)
6     **DECOUPLED_COMPUTE delivered on GCP + archived** (#270); **compute = GCP only**
(ADR-013, D0 dropped). **`COMPUTE_DECOUPLED_SERVING`** subproject spun up (#272, **ADR-014**) ?
one pure core + local/cloud profiles; **owner M0 design sign-off pending**. **Per-video
workspace P0** shipped (#264); **doc-status register + archival due-diligence** in place (#273).
The **khelsutra site + per-rally UI** ship in parallel (sibling track). The **rally-quality /
owned-model** track (below) is the core ML work, gated on **growing the golden corpus**.
7
8     ## ? Prioritized pending ? all tracks
9     **P0 ? highest leverage / unblocks the most**
10    0. *** [owner-priority, 2026-06-21] MAKE IT UI-DRIVEN + ALPHA-SHAREABLE ? "a few
recordings lying around ? highlights, from a page, no SSH".** Today the cloud flow is **SSH/CLI-
only** (`docs/CLOUD_GPU_DRYRUN_GUIDE.md`); **M12 does NOT fix it.** Needs (a separate khelsutra
`web/` SPA + deploy track, NOT an engine seam): **(i)** deploy the engine as a reachable
**service** (M7 ? Cloud Run service, ? owner GCP spend for a *persistent* endpoint), **(ii)**
the **SPA ingest/progress screen** (khelsutra `web/` B-P2: enter/upload video ? `api/process` ?
poll `api/jobs` ? reel; ZERO engine code), **(iii)** flip **M9 edge-auth ON** (merged #290,
```

default-OFF) to gate alpha tenants, ***(iv)*** a **compute PICKER** in the UI **"(your machine vs our cloud", D13/D15)**. The engine seams make it ***possible***; the SPA makes it ***usable***. Full record: **[[ui-driven-cloud-serving-gap]]**.

11 1. ? **COMPUTE_DECOUPLED_SERVING` M0?M9 LARGELY SHIPPED** (2026-06-21, ADR-014 ratified). M1?M4 (#279/#280/#282/#283), **M5 #286, M6 #287, M8 #288, M9-auth #290** all merged (default-OFF); the **real M7 L4 run** validated the worker (22 rallies, VM deleted, \$0). All \$8 gating Qs locked (D7?D17). **Owner residual** (not blocking Claude's default-OFF code): counsel ToS/DPA (M9-consent live + M11 live), real GCP prices+FX (M8 calibrate), Google OAuth app (M12 real), CF team-domain/AUD + ingress-lock; real cloud verification authorized within a **\$100/?10k cap** (D17).

12 2. **[owner, ongoing]** Grow the golden corpus ? multi-court badminton + ?3 matches/sport (TT first). The single highest-leverage move; the quality track ***and*** full-corpus training feed on it.

680. user (2026-06-21T21:48:27.466Z)

The file C:\Users\avidu\Projects\bhi-wt-docs\docs\README.md has been updated successfully. (file state is current in your context ? no need to Read it back)

681. user (2026-06-21T21:48:38.303Z)

The file C:\Users\avidu\Projects\bhi-wt-docs\docs\NEXT_STEPS.md has been updated successfully. (file state is current in your context ? no need to Read it back)

682. assistant (2026-06-21T21:48:43.594Z)

Committing the tracked doc + index pointers.

683. user (2026-06-21T21:48:55.293Z)

origin/master: de2bf14 feat(pkg): P2e ? lazy-torch so the LIGHT tier imports with NO torch installed (default-OFF) (#308)
=== status ===
M docs/NEXT_STEPS.md
M docs/README.md
?? docs/PACKAGING_PLAN.md
warning: in the working copy of 'docs/PACKAGING_PLAN.md', LF will be replaced by CRLF the next time Git touches it
173a305 docs(pkg): track the batteries-included app plan (PACKAGING_PLAN.md) + index pointers
To https://github.com/Khelsutra/badminton-highlight-indexer.git
* [new branch] docs/packaging-plan-tracked -> docs/packaging-plan-tracked
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

684. user (2026-06-21T21:49:06.459Z)

https://github.com/Khelsutra/badminton-highlight-indexer/pull/309
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer

685. assistant (2026-06-21T21:49:14.380Z)

PR #309 (docs) opened: [track PACKAGING_PLAN.md](https://github.com/Khelsutra/badminton-highlight-indexer/pull/309). While CI runs, let me write the next-session icebreaker handoff into memory and wire the index.

686. assistant (2026-06-21T21:49:41.977Z)

name: next-session-packaging-app
description: "START HERE for the batteries-included downloadable-app track. Icebreaker + ramp-up + the prioritized next steps (P3 GPU validation ? P2b-install ? P2c wizard). 11 PRs merged 2026-06-22."
metadata:
node_type: memory
type: project

? NEXT SESSION ? batteries-included downloadable app (icebreaker)

****You are here (2026-06-22):**** the ****LIGHT-tier engine is fully bundle-ready**** on ``origin/master``. A user can ``pip install`` the engine and run ****`indexer --app`**** ? it boots a local webserver and opens the ****bundled KhelSutra SPA**** single-origin, ****torch-free****, with resolved app-home state, an env-overridable ffmpeg, a DNS-rebinding-guarded localhost bind, single-instance locking, and safe DB upgrades. ****11 PRs merged**** this build (P0a?e, P1a/b, P2-launcher/P2a/P2d/P2e); ****0 open PRs; 0 GCP VMs (\$0).****

? Paste-ready icebreaker prompt for the new session

```
> Resume the **batteries-included downloadable-app** track. Ramp up: read
`docs/PACKAGING_PLAN.md` (the tracked doc ? $7 tracker is the source of truth + the "? NEXT
SESSION" block) and memory `[[packaging-app-plan]]`; `git pull --ff-only` first. Then work the
prioritized next steps **in order, one small PR each**, with the same diligence as the merged 11
(isolated `git worktree` off `origin/master`, explicit-path commits, full suite + lint +
coverage, adversarial review on logic PRs, watch CI, merge on green):
> 1. **P3 GPU validation + screencast** ? provision an L4 (`deploy/gcp/create_gpu_vm.sh`),
install **the wheel** on the fresh GPU box + the WASB env, validate the packaged stack end-to-
end (`indexer --app` ? bundled SPA ? native-WASB ? reel), **make a screencast**, then
`deploy/gcp/teardown.sh` (deletes) + confirm $0. Confirm the exact cmd + hourly cost BEFORE
spending; stay under the **$100/?10k cap**; teardown is priority #1 (`[[gcp-vm-always-
delete]]`). ? Known blocker: native-WASB returned **0 rallies** on a fresh build last time ?
reproducing the 22-rally result needs the engine session's exact tuned `worker infer
detector_impl=native` config.
> 2. **P2b-install** ? a `uv tool install` / generated cross-platform install script (a clean
**torch-free** LIGHT install now, thanks to P2e) + ship a truly-local default `config.json` + an
uninstall manifest.
> 3. **P2c first-run wizard** ? the 3-step wizard (Gemini key / compute / videos) in the
**khelsutra `web/` SPA** rendered from `GET /api/capabilities`. CROSS-REPO
(`avidullu/khelsutra`, gitignored `dist/`) ? coordinate like P1b, then refresh `backend/ui/` via
`scripts/bundle_ui.sh`.
> Then P3-cloud (Cloud Run dispatch + micromamba Dockerfile = P3a) and P4 (ONNX export to
dissolve the dual-torch env + training + in-app annotation = the WASB-complete **north star**).
```

Ramp-up kit (read before scoping)

```
- **`docs/PACKAGING_PLAN.md`** ? the tracked doc (committed #309). $7 = source of truth; $2
decisions-locked (L1?L10, incl. 01 cross-platform/no-native + 02 LIGHT-first/north-star); $8
open owner gates (03 WASB-C3 counsel, 04 hosted-vs-BYO-cloud, 05?09).
- Memory `[[packaging-app-plan]]` (full per-PR detail) . `[[gcp-vm-always-delete]]` .
`[[compute-stack-gcp-only]]` . `[[gotcha-shared-checkout-branch-collisions]]` . `[[lesson-
verify-engine-surface-before-scoping]]`.
- Code anchors: `backend/utils/{app_home,ffmpeg,single_instance,redact}.py`; `backend/main.py`
(`indexer --app`, `_resolve_bind_host`, `/api/version|capabilities|videos|reels`, bundled-UI
serving); `backend/ui/` (committed SPA snapshot) + `scripts/bundle_ui.sh`;
`backend/infrastructure/database.py` (`SCHEMA_VERSION`, downgrade guard, backup);
`deploy/gcp/*`.
```

Standing constraints / gotchas

```
- **Shared checkout:** sibling sessions flip the main checkout's branch + merge concurrently
(`origin/master` moved ~6x mid-build). Always branch off `origin/master`, re-verify HEAD before
commit, stage **explicit paths** (never `git add -A`), rebase if a sibling lands first. There is
a sibling `serving/exposure-safety-preflight` branch reworking startup.py/main.py exposure
checks ? flag on its PR that it must keep my P0c `EXPOSED_HOST_ALLOWLIST_LOOPBACK`.
- **Licensing guardrail** (MIT/Apache/BSD only, code+data+weights): the CI `license-scan` lane
enforces it for python runtime deps; WASB **C3 weight provenance** is still ship-blocking
(`gate_verdict ship=False`) ? a public build must NOT auto-fetch the weights until counsel
clears (03).
- **GPU spend:** owner-authorized within the **$100/?10k cap**; announce exact cmd + cost,
DELETE after, confirm $0.
```

Done ? don't redo (11 PRs)

P0a app-home #293 · P0b /api/config redaction #295 · P0c bind+DNS-rebinding guard #296 · P0d scipy + license-scan/built-wheel CI lanes #297 · P0e single-instance lock #298 · P1a endpoints #300 · P1b SPA bundled single-origin #302 · P2-launcher `indexer --app` #303 · P2a ffmpeg resolver #305 · P2d DB-upgrade safety #307 · P2e lazy-torch (torch-free LIGHT) #308 · docs-tracked #309.

687. user (2026-06-21T21:49:42.021Z)

File created successfully at: C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\next-session-packaging-app.md (file state is current in your context ? no need to Read it back)

688. assistant (2026-06-21T21:49:54.026Z)

Now wire it into the MEMORY.md index. Let me read the current "? Start here" + packaging section.

689. user (2026-06-21T21:49:54.846Z)

```
1      # Memory Index
2
3      ## ? Start here
4      - [? START HERE ? khelsutra UX/l10n/runnable-tool](next-session-ux-localization.md) ?
**refreshed 2026-06-22.** The runnable tool is **BUILT + fully localized (6 langs
en/kn/hi/es/da/id)**: in-UI upload (T4) ? process ? pick ? **reel watermarked in the user's
language** ? download, + PMF locale analytics. DONE this session: A1/A4/A5/A8/A9, T4,
MULTILINGUAL-A, B3(#70), T11(#299), **locale?watermark+analytics (LW1 #304 / LW2 #73 / LW3 #306,
verified via real watermark video for all 6 locales)**. **? ICEBREAKER ($0): ONE thing left to
make it alpha-real = stand up the T2 deploy** (owner: CF Access + GEMINI_API_KEY rotation +
cloudflared/CF-Pages per `deploy/DEPLOY_ALPHA.md`); THEN native-speaker review (launch gate) +
A3/A6/A7 + MULTILINGUAL Phase B/C (gated). *(Engine/serving track = own [session-
handoff](session-handoff.md); don't conflate.)*
5      - [Current state](current-state.md) ? **live micro-checkpoint** (updated every
checkpoint): open PRs, what just landed, next steps, open decisions. **Read this first.** Now
LIVE-only (handoff + today's checkpoints); kept read-first-short.
6      - [Current-state archive](current-state-archive.md) ? chronological checkpoint history
(2026-06-06 ? 06-18, 96 blocks) split out of current-state 2026-06-19. Deep history only; NOT
read each session.
7      - [Session handoff](session-handoff.md) ? current state · next steps · open decisions ·
ramp-up kit. AUTO-refreshed at meaningful checkpoints (PR merged, key decision, phase shift,
session wrap) ? see working-agreement-checkpoints staleness tripwire.
8      - [Code map artifact](code-map-artifact.md) ? for CODE/architecture ramp-up, read
`docs/CODE_MAP.md` FIRST (dense anchor-rich map; replaces a file crawl). How to regenerate it.
9      - [File-placement conventions](file-placement-conventions.md) ? **where new files go**
(scripts/ vs backend/eval/ vs tools/ vs scratch/ vs training/ vs artifacts/ ?); **read before
committing a new file.** Authoritative table = `docs/CODE_MAP.md` $0 (PR #138); scratch/
gitignored (#136).
10
```

690. user (2026-06-21T21:50:05.739Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

691. user (2026-06-21T21:50:11.628Z)

```
41:## Packaging / distribution
42:- [Batteries-included app plan](packaging-app-plan.md) ? **NEW owner track (2026-06-21):**
ship the engine as a downloadable batteries-included Python app (local web service for
inference/training/annotation, local+cloud). Tracked doc = `docs/PACKAGING_PLAN.md`; **01
(cross-platform Win/Mac/Linux, NO native component) + 02 (LIGHT-first, WASB-complete = sole
completion goal) LOCKED**; 03+ gated. Pairs with [ui-driven-cloud-serving-gap](ui-driven-cloud-
serving-gap.md). Don't relitigate the locked calls.
```

692. user (2026-06-21T21:50:22.157Z)

The file C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\memory\MEMORY.md has been updated successfully. (file state is current in your context ? no need to Read it back)

693. assistant (2026-06-21T21:50:27.292Z)

Handoffs are wired (MEMORY.md ? Start here + the icebreaker memory + the tracked doc + README/NEXT_STEPS). Let me confirm the docs PR CI and merge it.

694. user (2026-06-21T21:51:51.780Z)

```
Syntax sweep + import smoke (no GPU stack)      pass    21s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081065
Built wheel smoke (install + launch)      pass    1m14s  https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27918534788/job/82608081054
Dependency license scan (no copyleft in runtime deps) pass    31s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081204
Full test suite (offline, mocked)      pass    2m23s  https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27918534788/job/82608081044
Golden fixtures cross-OS (macos-latest) pass    20s  https://github.com/Khelsutra/badminton-
highlight-indexer/actions/runs/27918534788/job/82608081110
Golden fixtures cross-OS (ubuntu-latest)      pass    20s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081055
Golden fixtures cross-OS (windows-latest)      pass    1m6s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081079
Mypy (lenient baseline) pass    32s  https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081115
Ruff lint      pass    8s  https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081217
Syntax sweep + import smoke (no GPU stack)      pass    21s
https://github.com/Khelsutra/badminton-highlight-
indexer/actions/runs/27918534788/job/82608081065
Shell cwd was reset to C:\Users\avidu\Projects\badminton-highlight-indexer
```

695. assistant (2026-06-21T21:51:55.622Z)

All 9 lanes green. Merging the docs PR and cleaning up.

696. user (2026-06-21T21:52:05.194Z)

failed to run git: fatal: 'master' is already used by worktree at
'C:/Users/avidu/Projects/badminton-highlight-indexer/.claude/worktrees/tier1-windowing'

```
de2bf14..a42c132  master    -> origin/master
=== master ===
a42c132 docs(pkg): track the batteries-included app plan (PACKAGING_PLAN.md) + index pointers
(#309)
=== open PRs (mine) ===
309     docs(pkg): track the batteries-included app plan (PACKAGING_PLAN.md) + index pointers
docs/packaging-plan-tracked      MERGED  2026-06-21T21:49:05Z
=== PACKAGING_PLAN.md now on master? ===
YES ? on master
```

697. assistant (2026-06-21T21:52:35.745Z)

Everything's fresh and findable. Docs PR #309 merged ? **`docs/PACKAGING\_PLAN.md` is now on master** (12 PRs this session), indexed in `docs/README.md` + `docs/NEXT\_STEPS.md`, and the memory handoffs are wired (`? Start here ? next-session-packaging-app` + the full-detail

`packaging-app-plan`). 0 open PRs, 0 VMs (\$0).

? Icebreaker prompt ? paste this to start the new session

...

Resume the batteries-included downloadable-app track for the badminton-highlight-indexer.

Ramp up first: `git pull --ff-only`, then read docs/PACKAGING_PLAN.md (the tracked doc ? \$7 progress tracker is the source of truth, and the "? NEXT SESSION" block at the top) and the memory note next-session-packaging-app (under ? Start here). Don't relitigate the locked decisions (L1?L10) or redo the 11 merged PRs.

You are here: the LIGHT-tier engine is fully bundle-ready on master ? `pip install` ? `indexer --app` boots a local webserver and serves the bundled KhelSutra SPA single-origin, torch-free, with app-home state / env-overridable ffmpeg / DNS-rebinding-guarded localhost bind / single-instance lock / safe DB upgrades.

Work the remaining items IN ORDER, one small PR each, with the same diligence as the merged 11 (isolated git worktree off origin/master, explicit-path commits, full suite + ruff + mypy + coverage, adversarial review on logic PRs, watch CI, merge on green ? shared checkout, so re-verify HEAD before commit):

1. P3 GPU validation + screencast. Provision an L4 via deploy/gcp/create_gpu_vm.sh, install THE WHEEL on the fresh GPU box + the WASB env, validate the packaged stack end-to-end (indexer --app ? bundled SPA ? native-WASB ? reel), make a screencast, then deploy/gcp/teardown.sh (deletes) and confirm \$0. Announce the exact cmd + hourly cost BEFORE spending; stay under the \$100/?10k cap; teardown is priority #1. Known blocker: native-WASB returned 0 rallies on a fresh build last time ? reproducing the 22-rally result needs the engine session's exact tuned
`worker infer detector\_impl=native` config; budget time to nail it.
2. P2b-install ? a `uv tool install` / generated cross-platform install script (a clean torch-free LIGHT install now, thanks to P2e) + ship a truly-local default config.json + an uninstall manifest.
3. P2c first-run wizard ? the 3-step wizard (Gemini key / compute / videos) in the khelsutra web/ SPA, rendered from GET /api/capabilities. CROSS-REPO (avidullu/khelsutra, gitignored dist) ? coordinate like P1b, then refresh backend/ui/ via scripts/bundle_ui.sh.

Then P3-cloud (Cloud Run dispatch + micromamba Dockerfile = P3a) and P4 (ONNX export to dissolve the dual-torch env + BYO training + in-app annotation = the WASB-complete north star).

One coordination note: a sibling branch `serving/exposure-safety-preflight` reworks startup.py/main.py exposure checks ? when it opens a PR, make sure it keeps my P0c EXPOSED_HOST_ALLOWLIST_LOOPBACK check.

...

(The same prompt is saved verbatim in the `next-session-packaging-app` memory, so a new session will surface it during ramp-up even without the paste.)

What's fresh for the new session

- **`docs/PACKAGING\_PLAN.md`** (on master) ? `Status: IN PROGRESS`, \$7 tracker shows P0a?e / P1a/b / P2-launcher/P2a/P2d/P2e ? with PR links, and a top ***? NEXT SESSION ? start here*** block.
- **`docs/README.md`** / **`docs/NEXT\_STEPS.md`** ? point to it; `DESKTOP\_APP\_PLAN.md` marked superseded.
- **Memory:** `next-session-packaging-app` (icebreaker, under ? Start here) + `packaging-app-plan` (per-PR detail) + the MEMORY.md index line ? all refreshed; cross-linked to `[[gcp-vm-always-delete]]`, `[[gotcha-shared-checkout-branch-collisions]]`, `[[ui-driven-cloud-serving-gap]]`.

Thanks for the kind words ? this was a great run: ****12 PRs**** taking the engine from "git-clone-and-build-from-source" to a ****torch-free, bundle-ready, one-command app**** (``indexer --app``), with the GPU validation + install script + wizard cleanly teed up. Whenever you start the new session, that prompt drops you straight into P3.