

# claude-windows / f7103c4d-c406-4c6f-9a8a-f4d47f2e54eb

## Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\f7103c4d-c406-4c6f-9a8a-f4d47f2e54eb\subagents\workflows\wf_f5b294af-947\agent-a5382f6331480f689.jsonl`  
- SHA-256: `d153f1c8a54f19ef207934f6fe974477632a97fe99aa78859054f89deece3f33`  
- Source modified: `2026-06-22T18:13:02+00:00`  
- Imported at: `2026-07-05T16:48:22+00:00`  
- project: `wf_f5b294af-947`  
- session_id: `f7103c4d-c406-4c6f-9a8a-f4d47f2e54eb`
```

## Transcript

### 1. user (2026-06-22T18:11:00.478Z)

Research the CURRENT Google Gemini API pricing (it is mid-2026; prices and model aliases change, so VERIFY via the web ? do not answer from memory). Use WebSearch + WebFetch against the official pages (ai.google.dev/gemini-api/docs/pricing, ai.google.dev/gemini-api/docs/tokens, cloud.google.com Vertex pricing). TARGET: the model alias "gemini-flash-latest" ? first determine what it currently maps to (e.g. gemini-2.5-flash, or a newer flash), then report that model PAID-tier pricing: USD per 1M INPUT tokens and per 1M OUTPUT tokens. ALSO report how Gemini bills VIDEO input: tokens consumed per SECOND of video at DEFAULT media resolution, and at LOW media resolution (the cheaper setting), with the exact documented numbers. ALSO: is there a FREE tier for this model through the Gemini API (Google AI Studio key), and what are its limits ? because if so the actual incurred cost can be \$0 within limits. Return precise numbers with the source URLs and flag any uncertainty. Context for sanity-checking: segmenting a ~500-second 1080p video in ~9 chunks of ~60s each (downsized before upload), output is a small JSON list of rally time-ranges (a few hundred tokens per chunk). FOCUS your search on the model alias resolution + the per-1M token prices.

### 2. assistant (2026-06-22T18:11:03.764Z)

I'll research the current Gemini API pricing. Let me start by loading the web tools and fetching the official pricing pages.

### 3. assistant (2026-06-22T18:11:07.164Z)

Let me search for the current model alias resolution and pricing.

### 4. user (2026-06-22T18:11:15.669Z)

Web search results for query: "gemini-flash-latest model alias maps to which model 2026"

Links: [{"title": "Models | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/models"}, {"title": "Learn about supported models | Firebase AI Logic", "url": "https://firebase.google.com/docs/ai-logic/models"}, {"title": "?Gemini Apps? release updates & improvements", "url": "https://gemini.google/release-notes/"}, {"title": "Continuing to bring you our latest models, with an improved Gemini 2.5 Flash and Flash-Lite release - Google Developers Blog", "url": "https://developers.googleblog.com/en/continuing-to-bring-you-our-latest-models-with-an-improved-gemini-2-5-flash-and-flash-lite-release/"}, {"title": "Gemini 3 Flash | Gemini Enterprise Agent Platform | Google Cloud Documentation", "url": "https://docs.cloud.google.com/vertex-ai/generative-ai/docs/models/gemini/3-flash"}, {"title": "Release notes | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/changelog"}, {"title": "7 Gemini Models Compared: Which to Use in 2026 (Pro, Flash, Ultra)", "url": "https://teamai.com/blog/language-models-llms/gemini-models-explained-the-complete-2026-guide/"}, {"title": "Gemini (language model) - Wikipedia", "url": "https://en.wikipedia.org/wiki/Gemini\_(language\_model)"}], {"title": "Gemini 3.5:"

```
frontier intelligence with action", "url": "https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-3-5/"]]
```

Based on the search results, I found information about the `-latest` model alias system for Gemini models. Google introduced a `-latest` alias for each model family that always points to the most recent model versions, allowing developers to experiment with new features without needing to update code for each release.

However, the search results don't contain specific information about what exact model version `gemini-flash-latest` currently maps to as of June 2026. The search results mention experimental models like `gemini-2.5-flash` and newer versions like `gemini-3.5-flash`, but don't explicitly state the current mapping for the `gemini-flash-latest` alias.

To get the most current and accurate mapping, I recommend checking Google's official Gemini API documentation at <https://ai.google.dev/gemini-api/docs/models>, which was last updated on June 15, 2026 according to the search results.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

## 5. user (2026-06-22T18:11:18.684Z)

Web search results for query: "Gemini API pricing gemini-flash-latest input output tokens per 1M 2026"

Links: [{"title": "Gemini Developer API pricing | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/pricing"}, {"title": "Agent Platform Pricing | Google Cloud", "url": "https://cloud.google.com/gemini-enterprise-agent-platform/generative-ai/pricing"}, {"title": "Gemini API Pricing Calculator & Cost Guide (Jun 2026)", "url": "https://costgoat.com/pricing/gemini-api"}, {"title": "Gemini API Pricing May 2026: 3.5 Flash, 3.1 Pro, 2.5 Lite", "url": "https://www.metacto.com/blogs/the-true-cost-of-google-gemini-a-guide-to-api-pricing-and-integration"}, {"title": "Gemini Pricing in 2026 for Individuals, Orgs & Developers", "url": "https://www.finout.io/blog/gemini-pricing-in-2026"}, {"title": "Gemini API Pricing: Full Breakdown of Costs (Jun 2026)", "url": "https://developer.puter.com/tutorials/gemini-api-pricing/"}, {"title": "Gemini API Pricing 2026: Complete Per-1M-Token Cost Guide with Calculator - All 7 Models Compared | Free Tier Details | Cost Optimization Tips | AI Free API", "url": "https://www.aifreeapi.com/en/posts/gemini-api-pricing-2026"}, {"title": "Gemini API \$ per 1M tokens (May 2026) | Nicola Lazzari", "url": "https://nicolalazzari.ai/articles/gemini-api-pricing-explained-2026"}, {"title": "Gemini Pricing 2026: Plans, API & Workspace Cost Guide", "url": "https://felloai.com/gemini-pricing/"}, {"title": "Google Gemini API Pricing May 2026: Gemini 3.1 Pro, Free Tier Changes | AI Pricing Guru", "url": "https://www.aipricing.guru/google-ai-pricing/"}]

Based on the search results, here's the pricing information for the Gemini API in 2026:

## Gemini API Pricing for Input/Output Tokens per 1M

The latest model based on Gemini 2.5 Flash lite optimized for cost-efficiency, high throughput and high quality is referenced, but the search results don't contain pricing specifically labeled for "gemini-flash-latest."

However, here are the current Flash models and their pricing as of June 2026:

- Gemini 3.5 Flash launched on May 19, 2026 with a 1M-token input window and 64K output at a global rate of \$1.50 / \$9.00 per million tokens with \$0.15 cached input
- Gemini 3 Flash costs \$0.50 input / \$3.00 output per million tokens and retains free tier with reduced quota, while Gemini 3.1 Flash-Lite costs \$0.25 input / \$1.50 output per million tokens and also retains free tier with reduced quota
- Gemini 2.5 Flash-Lite is the lowest-cost option at \$0.10/1M input tokens and \$0.40/1M output tokens, with batch mode rates dropping to \$0.05 input / \$0.20 output

If you're looking for a specific "gemini-flash-latest" model ID, you may want to check Google's official pricing page directly, as the search results reference multiple Flash variants with different performance and cost profiles.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

#### 6. assistant (2026-06-22T18:11:21.634Z)

The search results show the model landscape has evolved significantly by mid-2026 (Gemini 3.x, 3.5 Flash exist). I need to verify the exact alias resolution and pricing from official pages. Let me fetch the official docs directly.

#### 7. user (2026-06-22T18:11:34.193Z)

## Gemini Flash Models Pricing Summary

Based on the pricing page content, here's what I found:

### gemini-flash-latest

This model name does not appear on the pricing page.

### Gemini 2.5 Flash

**\*\*Paid Tier (Standard):\*\*** "\$0.30 (text / image / video) \$1.00 (audio)" input; "\$2.50" output per 1M tokens

**\*\*Free Tier:\*\*** Yes - "Free of charge" for both input and output, with "500 RPD" for Google Search grounding (shared with Flash-Lite)

### Gemini 3 Flash Preview

**\*\*Paid Tier (Standard):\*\*** "\$0.50 (text / image / video) \$1.00 (audio)" input; "\$3.00" output per 1M tokens

**\*\*Free Tier:\*\*** Yes - "Free of charge" for input and output, with "5,000 prompts per month (free, shared across Gemini 3)" for Search grounding

### Gemini 3.5 Flash

**\*\*Paid Tier (Standard):\*\*** "\$1.50" input; "\$9.00" output per 1M tokens

**\*\*Free Tier:\*\*** Yes - "Free of charge" for both input and output, with "5,000 prompts per month (free, shared across Gemini 3)" for Search grounding

#### 8. user (2026-06-22T18:11:35.773Z)

## Latest Model Alias Mappings

Based on the documentation provided, **\*\*the specific model ID that "gemini-flash-latest" currently resolves to is not explicitly stated\*\*** in this content.

### What the documentation says about latest aliases:

The page explains that latest versions work this way: "Points to the latest release for a specific model variation. This can be a stable, preview or experimental release. This alias will get hot-swapped with every new release of a specific model variation."

The only "-latest" alias example given is:

> "For example: `gemini-flash-latest`."

However, this is presented merely as an example format, not as a statement of what it currently

maps to.

## Other alias patterns mentioned:

- **Stable aliases**: Point to specific stable models (example: ``gemini-3.5-flash``)
- **Preview aliases**: Point to preview versions (example: ``gemini-2.5-flash-preview-09-2025``)
- **Experimental aliases**: Point to experimental models

The documentation emphasizes that latest aliases change with releases and provides 2-week notice before updates, but does not provide a current resolution table for specific latest aliases.

### 9. assistant (2026-06-22T18:11:39.840Z)

The official pricing page lists multiple Flash models but doesn't show what ``gemini-flash-latest`` maps to. The models page confirms ``gemini-flash-latest`` exists as an alias but uses it as an example. I need to nail down the current resolution. Let me check the changelog and a couple of corroborating sources, plus the video token billing.

### 10. user (2026-06-22T18:11:46.617Z)

## Gemini Flash Latest Model Updates

Based on the release notes, here are the documented changes to the ``gemini-flash-latest`` alias:

**January 21, 2026:**

The changelog states: "gemini-flash-latest switched to gemini-3-flash-preview"

**September 24, 2024:**

An earlier update noted: "Updated the gemini-1.5-flash-latest model code to use gemini-1.5-flash-002"

**Most Recent:**

As of the latest entries in the documentation, ``gemini-flash-latest`` currently resolves to ``gemini-3-flash-preview``, which was the most recent hot-swap mentioned in the changelog. This represents the shift to the Gemini 3 series announced in January 2026.

### 11. user (2026-06-22T18:11:47.706Z)

## Token Consumption Per Second of Video

According to the documentation, Gemini's token consumption rates for video are:

**At default media resolution:** "Approximately 300 tokens per second of video"

**At low media resolution:** "100 tokens per second of video"

## How Token Count Is Calculated

The breakdown includes:

- **Frame sampling:** Video is processed at 1 frame per second (FPS)
- **Frame tokenization:** At low resolution, each frame uses 66 tokens; at default resolution, each frame consumes 258 tokens
- **Audio:** "32 tokens per second"
- **Metadata:** Also included in the total

The ``media_resolution`` parameter directly controls "the maximum number of tokens allocated per input image or video frame," with higher resolutions enabling better fine-text reading and small-detail identification at the cost of increased token usage and latency.

### 12. assistant (2026-06-22T18:11:51.404Z)

I now have the key facts. `gemini-flash-latest` resolved to `gemini-3-flash-preview` as of Jan 21, 2026. Let me confirm this is still current (no later hot-swap) and double-check the Gemini 3 Flash Preview pricing and free-tier limits with one more corroborating fetch.

### 13. user (2026-06-22T18:11:58.212Z)

Based on the provided content, I cannot find specific FREE tier rate limits (RPM, TPM, RPD) for Gemini 3.5 Flash or other Flash models.

The page states: "Rate limits depend on a variety of factors (such as your usage tier) and can be viewed in Google AI Studio."

It directs users to "[View your active rate limits in AI Studio](https://aistudio.google.com/rate-limit?timeRange=last-28-days)" rather than publishing specific numbers for the Free tier. The document only provides detailed Batch API rate limits for Tiers 1-3, with no equivalent Free tier batch specifications listed.

To find your actual Free tier limits, you would need to check your project settings directly in Google AI Studio.

### 14. user (2026-06-22T18:12:03.532Z)

Web search results for query: "\"gemini-flash-latest\" \"gemini-3-flash\" changelog hot-swap 2026"

Links: [{"title": "Release notes | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/changelog"}, {"title": "Gemini 3 Flash | Gemini Enterprise Agent Platform | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-flash"}, {"title": "?Gemini Apps? release updates & improvements", "url": "https://gemini.google/release-notes/"}, {"title": "Models | Gemini API | Google AI for Developers", "url": "https://ai.google.dev/gemini-api/docs/models"}, {"title": "Gemini 3.1 Flash-Lite | Gemini Enterprise Agent Platform | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-1-flash-lite"}, {"title": "This page documents updates to the Gemini API. May 28, 2026", "url": "https://ai.google.dev/gemini-api/docs/changelog.md.txt"}, {"title": "Introducing Gemini 3 Flash: Benchmarks, global availability", "url": "https://blog.google/products/gemini/gemini-3-flash/"}, {"title": "Gemini 3.2 Flash News: Official Status, Leak Signals, and What Developers Should Use Now | LaoZhang AI Blog", "url": "https://blog.laozhang.ai/en/posts/gemini-3-2-flash-news"}, {"title": "Gemini Code Assist release notes | Google for Developers", "url": "https://developers.google.com/gemini-code-assist/resources/release-notes"}, {"title": "Gemini 3.5 Flash | Gemini Enterprise Agent Platform | Google Cloud Documentation", "url": "https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-5-flash"}]

Based on the search results, here's what I found regarding your query about `gemini-flash-latest`, `gemini-3-flash`, changelog, and hot-swap in 2026:

## Hot-Swap Functionality

The `gemini-flash-latest` alias gets hot-swapped with every new release of a specific model variation, with a 2-week notice provided through email before the version behind latest is changed.

## Gemini 3 Flash Model

Gemini 3 Flash combines Gemini 3 Pro's reasoning capabilities with the Flash line's levels on latency, efficiency, and cost, and is designed to tackle complex agentic workflows. Gemini 3 Flash is now the new default model in the Gemini app and offers next generation intelligence at lightning speed, representing a major capability upgrade over Gemini 2.5 Flash.

## 2026 Changelog Highlights

Recent changelog entries from 2026 include:

- Released `gemini-3.5-flash`, the generally available (GA) version of Gemini 3.5 Flash on May 19, 2026, which is the most intelligent model for sustained frontier performance on agentic and coding tasks.
- Released `gemini-3.1-flash-lite` on May 7, 2026, the generally available (GA) version optimized for speed, scale, and cost efficiency.
- The `gemini-3.1-flash-lite-preview` model was shut down on May 25, 2026, with users advised to use `gemini-3.1-flash-lite` instead.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

#### 15. assistant (2026-06-22T18:12:07.379Z)

There may be a later hot-swap I should verify. The Jan 21, 2026 entry switched it to `gemini-3-flash-preview`, but GA `gemini-3.5-flash` (May 19) and `gemini-3.1-flash-lite` (May 7) have since shipped. Let me fetch the raw changelog text to check for any `gemini-flash-latest` re-pointing after January, and grab the rate-limits txt for free-tier numbers.

#### 16. user (2026-06-22T18:12:14.141Z)

### Gemini Flash Latest Alias Changes

Based on the documentation, here are all mentions of `gemini-flash-latest` redirects:

#### Change 1: November 21, 2024

**\*\*Switched to:\*\*** `gemini-1.5-flash-002`

Quote: "Updated the `gemini-1.5-flash-latest` and `gemini-1.5-flash` model aliases to use `gemini-1.5-flash-002`."

#### Change 2: January 21, 2026

**\*\*Switched to:\*\*** `gemini-3-flash-preview`

Quote: "`gemini-flash-latest` switched to `gemini-3-flash-preview`"

---

### Most Recent Update

**\*\*Date:\*\*** January 21, 2026

**\*\*Current pointer:\*\*** According to the documentation, `gemini-flash-latest` now directs to `gemini-3-flash-preview`, described as the model delivering "fast frontier-class performance that rivals larger models at a fraction of the cost."

#### 17. user (2026-06-22T18:12:14.664Z)

I cannot provide the specific Free tier rate limits (RPM, TPM, RPD) for those models because **\*\*the webpage content does not include a Free tier rate limits table\*\***.

The document states: "Rate limits depend on a variety of factors (such as your usage tier) and can be viewed in Google AI Studio," but directs users to check the actual rates in the AI Studio interface rather than listing them in this article.

The only rate limit tables presented are for **\*\*Batch API enqueued tokens\*\*** across Tier 1, Tier 2, and Tier 3?not standard interactive rate limits for the Free tier.

#### 18. assistant (2026-06-22T18:12:19.324Z)

Confirmed: the most recent hot-swap of `gemini-flash-latest` was January 21, 2026 ?  
`gemini-3-flash-preview`, with no later re-pointing in the changelog. So it currently resolves to **gemini-3-flash-preview**, priced at \$0.50 input / \$3.00 output per 1M tokens (text/image/video).

For the free-tier RPM/TPM/RPD numbers, the official rate-limits page now defers to AI Studio, but I have the pricing-page confirmation that Gemini 3 Flash Preview has a free tier. Let me get one corroborating source for the free-tier numeric limits.

## 19. user (2026-06-22T18:12:29.937Z)

Web search results for query: "Gemini 3 Flash preview free tier rate limits RPM RPD TPM Google AI Studio 2026"

Links: [{"title":"5 Ways to Solve AI Studio Gemini 3 Pro Rate Limits ? 2026 Complete Guide - Apiyi.com Blog","url":"https://help.apiyi.com/en/ai-studio-gemini-3-pro-rate-limit-solution-en.html"}, {"title":"Rate limits | Gemini API | Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/rate-limits"}, {"title":"Gemini API Free Tier 2026: Limits, Quotas, and What You Actually Get","url":"https://pecollective.com/tools/gemini-free-tier-guide/"}, {"title":"Rate limits - Gemini Flash 3.0 and 2.5 - Gemini Apps Community","url":"https://support.google.com/gemini/thread/400904312/rate-limits-gemini-flash-3-0-and-2-5?hl=en"}, {"title":"What are the RPD, RPM, and TPM limits for gemini-3.1-flash-tts-preview per tier? - Gemini API - Google AI Developers Forum","url":"https://discuss.ai.google.dev/t/what-are-the-rpd-rpm-and-tpm-limits-for-gemini-3-1-flash-tts-preview-per-tier/141793"}, {"title":"rate-limits.md.txt - Google AI for Developers","url":"https://ai.google.dev/gemini-api/docs/rate-limits.md.txt"}, {"title":"Gemini API Key Free in 2026: What Is Free, Where Limits Live, and Why More Keys Do Not Add Quota | LaoZhang AI Blog","url":"https://blog.laozhang.ai/en/posts/gemini-api-free-tier"}, {"title":"Interpreting Google AI Studio Rate Limits 2026 Latest Version: What to Do if Tier 1 RPD 250 is Too Strict - Apiyi.com Blog","url":"https://help.apiyi.com/en/google-ai-studio-rate-limits-2026-guide-en.html"}, {"title":"Gemini API Pricing Explained: Token Costs and Free Tier","url":"https://nicolalazzari.ai/articles/gemini-api-pricing-explained-2026"}]

I found relevant information about Gemini 3 Flash preview free tier rate limits in Google AI Studio. Here are the key findings:

## Gemini 3 Flash Preview Free Tier Rate Limits

For Gemini 3 Pro Preview, the free tier limits are roughly 10-50 RPM and 100+ RPD. However, there are important caveats to consider:

### **\*\*Important Context:\*\***

- Many users report that the actual limits feel much stricter than what's documented, which has been a common issue following quota adjustments in December 2025
- Some users have reported that the actual RPM is only about half of the documented value

### **\*\*TPM (Tokens Per Minute):\*\***

The search results don't provide specific TPM values for Gemini 3 Flash Preview on the free tier. Limits vary depending on the specific model being used, and some limits only apply to specific models.

### **\*\*General Information:\*\***

- Requests per day (RPD) quotas reset at midnight Pacific time
- Rate limits are more restricted for experimental and preview models
- Limits vary by tier and model ? check AI Studio for your project

For the most accurate and current rate limits for your specific project, it's recommended to check directly in Google AI Studio, as limits can vary by project, region, and account age.

REMINDER: You MUST include the sources above in your response to the user using markdown hyperlinks.

**20. assistant (2026-06-22T18:12:33.413Z)**

I have enough to answer with high confidence on the core facts (alias resolution, paid pricing, video token rates, free-tier existence), with the free-tier numeric RPM/RPD flagged as approximate since Google now surfaces exact numbers only in AI Studio. Let me record the structured output.

**21. user (2026-06-22T18:12:55.736Z)**

Structured output provided successfully

**22. assistant (2026-06-22T18:13:02.429Z)**

Done. Summary of verified findings:

- **gemini-flash-latest ? gemini-3-flash-preview** (official changelog, hot-swapped 2026-01-21; no later re-pointing despite GA 3.1-flash-lite and 3.5-flash shipping in May 2026).
- **Paid pricing (per 1M tokens):** \$0.50 input (text/image/video; audio \$1.00), \$3.00 output.
- **Video billing:** ~300 tokens/sec at default media resolution, ~100 tokens/sec at low (1 FPS sampling; 258 vs 66 tokens/frame + ~32 tokens/sec audio).
- **Free tier:** Yes via AI Studio key (input/output "Free of charge") ? so \$0 within limits; exact RPM/RPD numbers are no longer published in docs (deferred to AI Studio), flagged as a caveat.
- **Your use case:** a ~500s match ? \$0.03 (low res) to ~\$0.08 (default) on paid tier, or \$0 inside free-tier daily request limits; the binding constraint is RPD (9 requests/video), not token volume.

One flag worth surfacing: the alias points to a **preview** model, which can be deprecated on short notice and carries tighter free-tier limits ? pin `gemini-3-flash-preview` (or a stable Flash ID) explicitly if you need production stability.