

claude-windows / 7fce5410-4c57-4ec9-a054-be898bda22da

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7fce5410-4c57-4ec9-a054-be898bda22da\subagents\workflows\wf\_09101709-554\agent-a031a0970d39104d0.jsonl`
- SHA-256: `35d14a8783ca9963fab35082f44d4a431329c2c6a5912a17d57d605493c5b33c`
- Source modified: `2026-06-25T17:12:25+00:00`
- Imported at: `2026-07-05T16:48:26+00:00`
- project: `wf\_09101709-554`
- session_id: `7fce5410-4c57-4ec9-a054-be898bda22da`

Transcript

1. user (2026-06-25T17:11:13.831Z)

PR #381 ".github/workflows/nightly.yml" is CHECKED OUT (branch chore/nightly-quality-regression). It adds a Tier-1 nightly quality-regression GitHub Actions workflow (schedule cron "0 3 * * *" + workflow_dispatch) that re-runs the repo's OFFLINE guards (ruff, mypy, run_all_tests --cov-fail-under=70, golden-fixtures regression). The ONLY changed file is .github/workflows/nightly.yml (105 lines, ci.yml untouched). CRITICAL CONTEXT: a schedule/workflow_dispatch workflow does NOT self-trigger on the PR, so the PR's green CI only YAML-PARSED nightly.yml ? its STEPS were never executed. So you must verify the steps would actually run. Read BOTH .github/workflows/nightly.yml AND .github/workflows/ci.yml to compare. Ground every claim in the files.

Review ONLY scope/maintainability/conventions. (a) Single-purpose (1 file), Tier-1 only (NO gpu/real-video/run_regression-on-real-corpus)? (b) DRIFT RISK: nightly.yml appears to COPY ci.yml's job definitions ? as ci.yml evolves, the nightly silently diverges (a lane gets added/changed in ci.yml but not here). Would a reusable workflow (ci.yml exposing workflow_call, nightly calling it) be materially better, or is the duplication acceptable for 4 simple lanes? Give a clear recommendation, not just "consider". (c) Is the default-branch-only caveat documented (in the PR/workflow)? (d) name/clarity. Report findings + fix.

2. assistant (2026-06-25T17:11:16.210Z)

I'll review PR #381 by reading both workflow files and grounding every claim in their contents.

3. user (2026-06-25T17:11:16.712Z)

```
1     name: Nightly quality regression (Tier-1, offline)
2
3     # Tier-1 nightly quality guard. Per-PR `ci.yml` only fires on push/pull_request;
4     # this runs the SAME offline guards on a schedule so we also catch what per-PR CI
5     # can't: dependency / transitive-dep drift, flakiness over time, and rot during
6     # quiet periods when nothing is pushed.
7     #
8     # Tier-1 == OFFLINE ONLY: no GPU, no real video, no `run_regression.py` on the
9     # real corpus (that's Tier-2 ? owner-gated, needs a self-hosted GPU runner +
10    # the gitignored golden corpus). The lanes here mirror `ci.yml`'s offline jobs
11    # verbatim (same Python pin, same per-lane pip installs) so modules collect/run
12    # instead of silently skipping for a missing dep.
13    #
14    # NOTE: scheduled workflows only fire from the repo's DEFAULT branch, so this
15    # won't actually run on a cadence until it lands on `master`. Use the
16    # `workflow_dispatch` trigger to run it on demand for validation after merge.
17
18    on:
19      schedule:
20        # 03:00 UTC daily.
```

```

21 - cron: "0 3 * * *"
22 workflow_dispatch:
23
24 permissions:
25   contents: read
26
27 concurrency:
28   group: nightly- $\{\{ \text{github.ref} \}$ 
29   cancel-in-progress: true
30
31 jobs:
32   # Lint gate (ruff.toml at repo root) ? mirrors ci.yml::lint.
33   lint:
34     name: Ruff lint
35     runs-on: ubuntu-latest
36     steps:
37       - uses: actions/checkout@v6
38       - uses: actions/setup-python@v6
39         with:
40           python-version: "3.13"
41       - name: Install ruff
42         run: python -m pip install ruff
43       - name: Ruff check
44         run: ruff check . --output-format github
45
46   # Type gate (mypy.ini at repo root) ? mirrors ci.yml::typecheck. Typed core
47   # deps are installed so mypy sees real signatures (without them
48   # ignore_missing_imports would Any-stub half the checks).
49   typecheck:
50     name: Mypy (lenient baseline)
51     runs-on: ubuntu-latest
52     steps:
53       - uses: actions/checkout@v6
54       - uses: actions/setup-python@v6
55         with:
56           python-version: "3.13"
57       - name: Install mypy + typed core deps
58         run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
59 genai numpy
60       - name: Mypy
61         run: mypy backend training
62
63   # Full offline suite + coverage floor ? mirrors ci.yml::full-suite. CPU-only
64   # torch (no CUDA wheels on a CPU runner), then the editable install pulls the
65   # rest. Floor stays at the ci.yml value (70); never lower it without an owner
66   # decision.
67   full-suite:
68     name: Full test suite (offline, mocked)
69     runs-on: ubuntu-latest
70     steps:
71       - uses: actions/checkout@v6
72       - uses: actions/setup-python@v6
73         with:
74           python-version: "3.13"
75           cache: pip
76       - name: System libs for opencv-python
77         run: |
78           sudo apt-get update
79           sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
80 libgl1 libglu1-mesa-dev
81       - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
82         run: python -m pip install torch torchvision --index-url
83 https://download.pytorch.org/whl/cpu
84       - name: Install package (editable) + dev tooling
85         run: python -m pip install -e .[dev]

```

```

83         - name: Run full suite (with coverage floor)
84         run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
85
86         # Offline quality guard: the committed golden regression fixtures
87         # (docs/GOLDEN_REGRESSION_FIXTURES.md) ? synthetic + de-attributed real
88         # trajectory fixtures, NO video / GPU / DB. Mirrors the offline replay in
89         # ci.yml::golden-crossos (single-OS here; ci.yml already covers cross-OS on
90         # every PR ? nightly is about drift-over-time on the default branch). The
91         # CONTROL replay needs only numpy.
92         golden-fixtures:
93             name: Golden fixtures regression (offline quality guard)
94             runs-on: ubuntu-latest
95             steps:
96                 - uses: actions/checkout@v6
97                 - uses: actions/setup-python@v6
98                   with:
99                       python-version: "3.13"
100                 - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
101                   run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
102                 - name: Golden-fixtures characterization (offline)
103                   run: |
104                       python -m backend.eval.golden_fixtures --check
105                       python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
106

```

4. user (2026-06-25T17:11:17.200Z)

```

1     name: CI
2
3     on:
4         push:
5             branches: [master]
6         pull_request:
7
8     permissions:
9         contents: read
10
11     concurrency:
12         group: ci-${{ github.ref }}
13         cancel-in-progress: true
14
15     jobs:
16         # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17         # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
18         # tracked .py, including scripts/ and training/, which the backend-only
19         # py_compile smoke sweep does not reach.
20         lint:
21             name: Ruff lint
22             runs-on: ubuntu-latest
23             steps:
24                 - uses: actions/checkout@v6
25                 - uses: actions/setup-python@v6
26                   with:
27                       python-version: "3.13"
28                 - name: Install ruff
29                   run: python -m pip install ruff
30                 - name: Ruff check
31                   run: ruff check . --output-format github
32
33         # PII / attribution leak guard (tools/leak_scan.py): fails on a committed 32-hex MD5
34         # (the video_id shape) or a personal absolute user path in the shipped-code surface
35         # (backend/training/scripts/tools/eval_baselines; tests/ excluded ? legit fixtures).

```

```

Real
36     # names are scanned only via a gitignored private list locally; CI runs the structural
guard.
37     leak-scan:
38         name: Leak scan (PII / attribution guard)
39         runs-on: ubuntu-latest
40         steps:
41             - uses: actions/checkout@v6
42             - uses: actions/setup-python@v6
43               with:
44                 python-version: "3.13"
45             - name: Leak scan
46               run: python -m tools.leak_scan
47
48     # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
49     # per-module quarantine list ? ratchet by removing entries as modules clean
50     # up. Typed core deps are installed so mypy sees their real signatures
51     # (without them ignore_missing_imports would Any-stub half the checks).
52     typecheck:
53         name: Mypy (lenient baseline)
54         runs-on: ubuntu-latest
55         steps:
56             - uses: actions/checkout@v6
57             - uses: actions/setup-python@v6
58               with:
59                 python-version: "3.13"
60             - name: Install mypy + typed core deps
61               run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
genai numpy
62             - name: Mypy
63               run: mypy backend training
64
65     # Fast guard against the 7fbb0 class of breakage: a SyntaxError or
66     # module-level import error shipping while the suite stays green.
67     # Deliberately does NOT install cv2/torch/numpy ? the smoke tests must hold
68     # on machines without the GPU stack (test_import_smoke stubs what's absent).
69     import-smoke:
70         name: Syntax sweep + import smoke (no GPU stack)
71         runs-on: ubuntu-latest
72         steps:
73             - uses: actions/checkout@v6
74             - uses: actions/setup-python@v6
75               with:
76                 python-version: "3.13"
77             - name: Install minimal deps
78               run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai
79             - name: Run import smoke tests
80               run: python -m pytest tests/test_import_smoke.py -v
81
82     full-suite:
83         name: Full test suite (offline, mocked)
84         runs-on: ubuntu-latest
85         steps:
86             - uses: actions/checkout@v6
87             - uses: actions/setup-python@v6
88               with:
89                 python-version: "3.13"
90             - cache: pip
91             - name: System libs for opencv-python
92               run: |
93                 sudo apt-get update
94                 sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
95             - name: Install CPU-only torch (no CUDA wheels on a CPU runner)

```

```

96         # torch/torchvision are NOT in pyproject deps ? their CPU/CUDA wheels
97         # need an out-of-band index that PEP 621 can't express. Keep installing
98         # them separately here, before the editable install pulls the rest.
99         run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
100     - name: Install package (editable) + dev tooling
101       # Editable PEP 621 install: runtime deps + the `dev` group
102       # (pytest, pytest-cov, httpx for the fastapi TestClient, ruff, mypy).
103       run: python -m pip install -e .[dev]
104     # obsreport (private, soft dep) is absent here: reporting tests skip via
105     # pytest.importorskip ? run_all_tests.py prints which. The coverage floor
106     # is therefore calibrated against THIS lane's number (local runs with
107     # obsreport installed measure a little higher).
108     # Floor calibrated 2026-06-11 against this lane's measured 72% (local
109     # with obsreport: 73%). Ratchet it UP as coverage grows; never down
110     # without an owner decision.
111     - name: Run full suite (with coverage floor)
112       run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
113
114     # Cross-OS determinism guard for the committed golden regression fixtures
115     # (docs/GOLDEN_REGRESSION_FIXTURES.md). The CONTROL replay path is pure-stdlib
116     # (only numpy, no cv2/torch), so this is a FAST, focused job ? it does NOT run
117     # the full suite x3 OSes. It replays the Windows/Linux/macOS runner against the
118     # committed (parse-compared) snapshots, catching any cross-OS float drift beyond
119     # the 1e-6 tolerance before it can flake a contributor's PR on another machine.
120     golden-crossos:
121       name: Golden fixtures cross-OS (${{ matrix.os }})
122       strategy:
123         fail-fast: false
124         matrix:
125           os: [ubuntu-latest, windows-latest, macos-latest]
126       runs-on: ${{ matrix.os }}
127       steps:
128         - uses: actions/checkout@v6
129         - uses: actions/setup-python@v6
130         with:
131           python-version: "3.13"
132         - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
133           run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
134         - name: Golden-fixtures characterization (cross-OS determinism guard)
135           # Also runs the freeze-real / time-origin-reset / M2 refusal-gate tests
(synthetic /
136           # tmp_path only ? no GPU/network) so the reset metric-invariance (fix #2) and
the
137           # de-attribution paths get cross-OS coverage and surface any win/mac float
drift.
138           run: |
139             python -m backend.eval.golden_fixtures --check
140             python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
141
142     # Commercial-clean licensing guardrail (MIT/Apache/BSD only ? PACKAGING_PLAN.md P0d).
Scans the
143     # RUNTIME dependency closure (`pip install .` ? no torch/dev) and fails on any
copyleft
144     # (GPL/AGPL/LGPL/SSPL). torch (BSD-3) is installed out-of-band and is not in this
closure. This
145     # automates a guardrail that was human-only; the report is printed for review either
way.
146     license-scan:
147       name: Dependency license scan (no copyleft in runtime deps)
148       runs-on: ubuntu-latest
149       steps:

```

```

150         - uses: actions/checkout@v6
151         - uses: actions/setup-python@v6
152         with:
153             python-version: "3.13"
154         - name: Install runtime deps + pip-licenses
155           run: |
156             python -m pip install .
157             python -m pip install pip-licenses
158         - name: Report dependency licenses
159           run: python -m piplicenses --format=markdown --with-urls --order=license
160         - name: Fail on any copyleft (GPL/AGPL/LGPL/SSPL) in the runtime closure
161           run: |
162             if python -m piplicenses --format=plain | grep -iE "GPL|SSPL"; then
163                 echo "::error::A runtime dependency carries a copyleft (GPL/AGPL/LGPL/SSPL)
license ? the guardrail is MIT/Apache/BSD only. Review the report above."
164             exit 1
165           fi
166           echo "OK: no GPL/AGPL/LGPL/SSPL in the runtime dependency closure."
167
168     # Built-wheel smoke (PACKAGING_PLAN.md P0d): build the wheel, install it NON-editable
from a
169     # non-repo cwd, and launch the console script. Catches packaging regressions the
source suite
170     # can't ? namespace-package resolution (backend has no __init__.py), missing data
files, hidden
171     # imports. (Validated locally: a clean-venv wheel install + CPU torch runs `indexer
--help`.)
172     build-artifact:
173       name: Built wheel smoke (install + launch)
174       runs-on: ubuntu-latest
175       steps:
176         - uses: actions/checkout@v6
177         - uses: actions/setup-python@v6
178         with:
179             python-version: "3.13"
180         - name: System libs for opencv-python
181           run: |
182             sudo apt-get update
183             sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
184         - name: Build the wheel
185           run: |
186             python -m pip install build
187             python -m build --wheel
188         - name: Install the built wheel + CPU torch (out-of-band index)
189           run: |
190             python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
191             python -m pip install dist/*.whl
192         - name: Launch the installed console script from a non-repo dir
193           run: |
194             cd /tmp
195             indexer --help
196             python -c "import backend, training, backend.utils.app_home; print('wheel
import OK')"
197

```

5. assistant (2026-06-25T17:11:20.925Z)

Both files are fully read. Let me verify a few grounding facts before finalizing: whether `run\_all\_tests.py` accepts the exact flags used, whether the golden-fixtures module/entrypoints exist, and whether there's any PR/doc description carrying the default-branch caveat.

6. user (2026-06-25T17:11:23.100Z)

```
1a400c0 ci: add Tier-1 nightly quality-regression workflow
86705c1 chore: ruff format gemini_refine.py + trajectory_hybrid.py (#379)
b72ac85 fix(B2/P17): surface Gemini provider failures in all callers (#378)
71a8241 chore(repo): fix corrupted .gitignore (UTF-16 tail) + ignore artifacts/, track product
overview (#377)
f7fb36e docs(audit): fix two accuracy nits ? D2 DoD wording + main.py line count (#376)
---DIFF STAT vs master---
.github/workflows/nightly.yml | 105 ++++++++++++++++++++++++++++++++++++++
1 file changed, 105 insertions(+)
```

7. user (2026-06-25T17:11:23.435Z)

```
backend\eval\golden_fixtures.py
```

8. user (2026-06-25T17:11:23.529Z)

```
tests\test_golden_fixtures.py
tests\test_golden_fixtures_split.py
tests\test_golden_real_fixtures.py
```

9. assistant (2026-06-25T17:11:27.154Z)

Diff confirms exactly one changed file (105 lines), and the referenced module/test entrypoints all exist. Let me verify the `--check` flag on the golden_fixtures module and the `run\_all\_tests.py` arg passthrough, plus confirm the PR body for the default-branch caveat.

10. user (2026-06-25T17:11:29.177Z)

```
735: g.add_argument("--check", action="store_true",
```

11. user (2026-06-25T17:11:30.573Z)

```
10: python run_all_tests.py --cov=backend --cov-report=term-missing # coverage
37: args = list(sys.argv[1:])
```

12. user (2026-06-25T17:11:32.185Z)

```
{"body": "## What\n\nAdds `.github/workflows/nightly.yml` ? a **Tier-1 nightly quality-  
regression** workflow that runs the repo's **offline** quality guards on a schedule.\n\n-  
**Triggers:** `schedule: cron \"0 3 * * *\" (03:00 UTC daily) **and** `workflow_dispatch:` (on-  
demand).\n-  
**Lanes** (mirror `ci.yml`'s offline jobs verbatim ? same Python 3.13 pin + same  
per-lane pip installs so modules actually collect/run instead of skipping for a missing dep):\n-  
`ruff check . --output-format github`\n-  
`mypy backend training`\n-  
`python run_all_tests.py --cov=backend --cov=training --cov-report=term --cov-fail-under=70`\n-  
golden-fixtures regression ? `python -m backend.eval.golden_fixtures --check` + `pytest  
tests/test_golden_fixtures.py tests/test_golden_real_fixtures.py -q` (synthetic + de-attributed  
real **trajectory** fixtures; **no video / GPU / DB**)\n\n## Why\n\nThe quality/regression  
machinery exists but nothing runs on a cadence ? per-PR `ci.yml` only fires on  
`push`/`pull_request`. A nightly run of the same offline guards catches what per-PR CI can't:  
**dependency / transitive-dep drift, flakiness over time, and rot during quiet periods** when  
nothing is pushed.\n\n## Scope ? Tier-1 (offline) only\n\n-  
**No GPU, no real video, no  
`run_regression.py` on the real corpus.** That's **Tier-2** (real-video F1) ? owner-gated, needs  
a self-hosted GPU runner + the gitignored golden corpus, and is explicitly out of scope here.\n-  
Self-contained: **`ci.yml` is untouched**, keeping this a single-purpose change.\n\n## Caveats /  
how to validate\n\n-  
?? **Scheduled workflows only fire from the repo's default branch.** This  
nightly will **not** run on a cadence until it merges to `master`. After merge, validate it on  
demand via the **`workflow_dispatch`** trigger (Actions ? `\"Nightly quality regression (Tier-1,  
offline)\"` ? Run workflow).\n-  
This PR's own checks are the normal `ci.yml` `pull_request` run  
(the new file is just YAML; it doesn't self-trigger on the PR).\n\n## Follow-ups (NOT in this  
PR)\n\n-  
**M4 quality-floor** ? add a per-slug F1 floor in  
`eval_baselines/regression_baseline.json` + a windowing fingerprint (CONTROL-only) so the  
golden-fixtures regression asserts an actual **F1 floor** rather than just exercising the flow.  
See `docs/NEXT_STEPS.md` / `docs/GOLDEN_REGRESSION_FIXTURES.md`.\n-  
**Tier-2** ? real-video F1  
via a self-hosted GPU runner is a separate, **owner-gated** effort.\n\n? Generated with [Claude
```

```
Code](https://claude.com/claude-code)\n","headRefName":"chore/nightly-quality-regression","title":"ci: add Tier-1 nightly quality-regression workflow"}
```

13. assistant (2026-06-25T17:11:39.612Z)

Now I have everything grounded. Let me verify the `run\_all\_tests.py` argv passthrough handles `--cov-fail-under` (it builds args from `sys.argv[1:]` and forwards to pytest), and confirm the lane comparison against ci.yml is complete.

Key verification facts established:

- Diff is exactly 1 file, 105 insertions (matches "single-purpose").
- `nightly.yml` `lint`, `typecheck`, `full-suite` lanes are byte-for-byte identical to `ci.yml`'s versions (same Python 3.13, same pip installs, same commands).
- `run\_all\_tests.py` forwards `sys.argv[1:]` to pytest, so `--cov=backend --cov=training --cov-report=term --cov-fail-under=70` would actually run.
- `backend/eval/golden\_fixtures.py` has `--check`; both test files exist ? golden-fixtures lane would run.
- The default-branch caveat IS documented in both the workflow header (lines 14-16) and the PR body.
- Tier-1-only scope holds: no GPU, no real video, no `run\_regression.py` on corpus.

The central concern is the drift risk from copying ci.yml's job definitions. Let me quantify exactly which lanes are duplicated.

14. user (2026-06-25T17:11:40.848Z)

```
30         try:
31             return importlib.util.find_spec(module) is not None
32         except (ImportError, ValueError):
33             return False
34
35
36     def main() -> int:
37         args = list(sys.argv[1:])
38         # Default to quiet unless the caller chose a verbosity.
39         if not any(a in ("-q", "-v", "-vv", "-vvv", "--quiet", "--verbose") for a in args):
40             args = ["-q", *args]
41         cmd = [sys.executable, "-m", "pytest", *args]
42
43         obs = "available" if _installed("obsreport") else "NOT installed ? reporting tests
will SKIP"
44         bar = "=" * 72
45         print(bar)
46         print(" Full test suite: " + " ".join(cmd))
47         print(f" obsreport (optional reporting dep): {obs}")
48         print(bar, flush=True)
49
50         rc = subprocess.run(cmd).returncode
51
52         print(bar)
53         print(" RESULT: " + ("PASS - all tests green" if rc == 0 else f"FAIL - pytest exit
{rc}"))
54         print(bar)
55         return rc
56
57
58     if __name__ == "__main__":
59         raise SystemExit(main())
```

15. assistant (2026-06-25T17:11:46.516Z)

Confirmed: `run\_all\_tests.py` forwards all args verbatim to pytest, so the nightly's `full-suite` command would execute correctly.

I now have everything grounded across both files plus the supporting code. Compiling findings.

The lane comparison vs `ci.yml`:

- ci.yml has 8 jobs: `lint`, `leak-scan`, `typecheck`, `import-smoke`, `full-suite`, `golden-crossos`, `license-scan`, `build-artifact`.
- nightly.yml has 4 jobs: `lint`, `typecheck`, `full-suite`, `golden-fixtures` (single-OS subset of `golden-crossos`).
- `lint`, `typecheck`, `full-suite` are duplicated verbatim. This is the drift surface.

16. user (2026-06-25T17:12:13.901Z)

Structured output provided successfully

17. assistant (2026-06-25T17:12:25.858Z)

Review complete. Summary of findings for PR #381 (`.github/workflows/nightly.yml`):

****Scope/conventions all pass**** (grounded in the files):

- Single-purpose: exactly 1 file, 105 insertions, `ci.yml` untouched (confirmed via `git diff --stat`).
- Tier-1/offline-only: no GPU, no real video, no `run\_regression.py` on the corpus; golden lane runs only the existing `backend.eval.golden\_fixtures --check` (`--check` exists at `golden\_fixtures.py:735`) + two present test files.
- Default-branch caveat documented in ****both**** the workflow header (lines 14-16) and the PR body.
- Steps would actually run: `run\_all\_tests.py:37-41` forwards `sys.argv[1:]` straight to pytest, so the cov flags reach pytest-cov; `workflow\_dispatch` lets the owner exercise them on demand after merge.

****The one material issue (MEDIUM) ? drift risk, with a clear recommendation, not "consider":**** `lint`, `typecheck`, and `full-suite` in `nightly.yml` are byte-for-byte copies of `ci.yml`'s versions, and the file *itself* advertises* the coupling ("mirror ci.yml's offline jobs verbatim", floor "stays at the ci.yml value (70)"). Since the whole point of the nightly is to catch drift, the copy is the thing most likely to drift. ****Recommendation:** extract the shared offline lanes into a reusable `workflow\_call` file that both `ci.yml` and `nightly.yml` invoke via `uses:`. ****** That single-sources the lane set, the Python pin, and the coverage floor (a LOW finding flags the floor literal as a second silent-drift vector ? `ci.yml:110` explicitly anticipates ratcheting it up). Keep `ci.yml`'s PR-only jobs (leak-scan, import-smoke, license-scan, build-artifact, cross-OS matrix) out of the reusable file so nightly stays the Tier-1 subset.

If the owner wants to keep this PR to one file, the fallback is to ship as-is and file a tracked follow-up for the `workflow\_call` refactor.