

claude-windows / 7fce5410-4c57-4ec9-a054-be898bda22da

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7fce5410-4c57-4ec9-a054-be898bda22da\subagents\workflows\wf\_09101709-554\agent-a71d52f73ce035295.jsonl`
- SHA-256: `730b7809f9fdbb906fe7f11e0b0307545f2cfd49c822719c648f645a22245fb7`
- Source modified: `2026-06-25T17:13:34+00:00`
- Imported at: `2026-07-05T16:48:26+00:00`
- project: `wf\_09101709-554`
- session_id: `7fce5410-4c57-4ec9-a054-be898bda22da`

Transcript

1. user (2026-06-25T17:13:08.368Z)

Adversarially verify this PR-review finding for PR #381 (branch chore/nightly-quality-regression checked out; read .github/workflows/nightly.yml + ci.yml). Be skeptical ? default to INVALID if wrong, already-handled, or a non-issue.

Finding [MEDIUM] No timeout-minutes on any job ? a hung run inherits the 6-hour default and burns Actions minutes unattended

Location: .github/workflows/nightly.yml ? all jobs (lint, typecheck, full-suite, golden-fixtures)

Detail: None of the jobs set `timeout-minutes`, so each falls back to GitHub Actions' default 6-hour job timeout. A stalled step (pip dependency-resolver backtracking, an apt-get hang in the full-suite libgl1 install, a deadlocked test) on an unattended nightly can burn up to 6h of runner minutes before being killed ? and because of the failure-visibility gap above, nobody would notice. ci.yml has the same omission, but a scheduled unattended job is exactly where a silent hang is most costly, so it matters more here.

Suggestion: Applied: added `timeout-minutes` per job (lint 10, typecheck 10, full-suite 30, golden-fixtures 15, notify-on-failure 5) tuned to each lane's offline workload. Tune to observed durations. I also spawned a follow-up task to add matching timeouts to ci.yml's 7 jobs so the two files stay consistent.

Return isReal + adjustedSeverity (INVALID if not real) + one-paragraph reasoning grounded in the actual files.

2. assistant (2026-06-25T17:13:10.592Z)

I'll verify this finding by reading the actual workflow files.

3. user (2026-06-25T17:13:11.692Z)

```
1      name: Nightly quality regression (Tier-1, offline)
2
3      # Tier-1 nightly quality guard. Per-PR `ci.yml` only fires on push/pull_request;
4      # this runs the SAME offline guards on a schedule so we also catch what per-PR CI
5      # can't: dependency / transitive-dep drift, flakiness over time, and rot during
6      # quiet periods when nothing is pushed.
7      #
8      # Tier-1 == OFFLINE ONLY: no GPU, no real video, no `run_regression.py` on the
9      # real corpus (that's Tier-2 ? owner-gated, needs a self-hosted GPU runner +
10     # the gitignored golden corpus). The lanes here mirror `ci.yml`'s offline jobs
11     # verbatim (same Python pin, same per-lane pip installs) so modules collect/run
12     # instead of silently skipping for a missing dep.
13     #
14     # NOTE: scheduled workflows only fire from the repo's DEFAULT branch, so this
15     # won't actually run on a cadence until it lands on `master`. Use the
16     # `workflow_dispatch` trigger to run it on demand for validation after merge.
```

```

17
18     on:
19         schedule:
20             # 03:00 UTC daily.
21             - cron: "0 3 * * *"
22         workflow_dispatch:
23
24     permissions:
25         contents: read
26
27     concurrency:
28         group: nightly-${{ github.ref }}
29         cancel-in-progress: true
30
31     jobs:
32         # Lint gate (ruff.toml at repo root) ? mirrors ci.yml::lint.
33         lint:
34             name: Ruff lint
35             runs-on: ubuntu-latest
36             timeout-minutes: 10
37             steps:
38                 - uses: actions/checkout@v6
39                 - uses: actions/setup-python@v6
40                   with:
41                       python-version: "3.13"
42                 - name: Install ruff
43                   run: python -m pip install ruff
44                 - name: Ruff check
45                   run: ruff check . --output-format github
46
47         # Type gate (mypy.ini at repo root) ? mirrors ci.yml::typecheck. Typed core
48         # deps are installed so mypy sees real signatures (without them
49         # ignore_missing_imports would Any-stub half the checks).
50         typecheck:
51             name: Mypy (lenient baseline)
52             runs-on: ubuntu-latest
53             timeout-minutes: 10
54             steps:
55                 - uses: actions/checkout@v6
56                 - uses: actions/setup-python@v6
57                   with:
58                       python-version: "3.13"
59                 - name: Install mypy + typed core deps
60                   run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
61 genai numpy
62                 - name: Mypy
63                   run: mypy backend training
64
65         # Full offline suite + coverage floor ? mirrors ci.yml::full-suite. CPU-only
66         # torch (no CUDA wheels on a CPU runner), then the editable install pulls the
67         # rest. Floor stays at the ci.yml value (70); never lower it without an owner
68         # decision.
69         full-suite:
70             name: Full test suite (offline, mocked)
71             runs-on: ubuntu-latest
72             timeout-minutes: 30
73             steps:
74                 - uses: actions/checkout@v6
75                 - uses: actions/setup-python@v6
76                   with:
77                       python-version: "3.13"
78                       cache: pip
79                 - name: System libs for opencv-python
80                   run: |
81                       sudo apt-get update

```

```

81         sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
82     - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
83     run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
84     - name: Install package (editable) + dev tooling
85     run: python -m pip install -e .[dev]
86     - name: Run full suite (with coverage floor)
87     run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
88
89     # Offline quality guard: the committed golden regression fixtures
90     # (docs/GOLDEN_REGRESSION_FIXTURES.md) ? synthetic + de-attributed real
91     # trajectory fixtures, NO video / GPU / DB. Mirrors the offline replay in
92     # ci.yml::golden-crossos (single-OS here; ci.yml already covers cross-OS on
93     # every PR ? nightly is about drift-over-time on the default branch). The
94     # CONTROL replay needs only numpy.
95     golden-fixtures:
96       name: Golden fixtures regression (offline quality guard)
97       runs-on: ubuntu-latest
98       timeout-minutes: 15
99       steps:
100         - uses: actions/checkout@v6
101         - uses: actions/setup-python@v6
102           with:
103             python-version: "3.13"
104         - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
105         run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
106         - name: Golden-fixtures characterization (offline)
107         run: |
108           python -m backend.eval.golden_fixtures --check
109           python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
110
111     # Failure visibility ? a silent nightly is useless. Scheduled runs have no PR
112     # to redden and nobody is watching the Actions tab on a quiet day (the exact
113     # "rot during quiet periods" this workflow exists to catch). On ANY guard
114     # failure, open/refresh a single tracking issue so a human is pinged. Uses the
115     # built-in GITHUB_TOKEN (no secrets); `issues: write` is scoped to THIS job
116     # only ? the guard jobs above keep `contents: read`.
117     notify-on-failure:
118       name: Surface nightly failure (open/refresh tracking issue)
119       needs: [lint, typecheck, full-suite, golden-fixtures]
120       if: failure() && github.event_name == 'schedule'
121       runs-on: ubuntu-latest
122       timeout-minutes: 5
123       permissions:
124         contents: read
125         issues: write
126       steps:
127         - uses: actions/github-script@v7
128           with:
129             script: |
130               const runUrl =
`${context.serverUrl}/${context.repo.owner}/${context.repo.repo}/actions/runs/${context.runId}`;
131               const title = 'Nightly quality regression FAILED';
132               const body = [
133                 'The Tier-1 nightly offline quality guard failed.',
134                 '',
135                 ` - Run: ${runUrl}`,
136                 ` - Date (UTC): ${new Date().toISOString()}`,
137                 '',
138                 'One of ruff / mypy / full-suite (cov>=70) / golden-fixtures regressed.',
139                 'Open the run for the failing lane. Close this issue once green.',

```

```

140         ].join('\n');
141         // Reuse a single open tracking issue instead of spamming one per night.
142         const existing = await github.rest.issues.listForRepo({
143             owner: context.repo.owner,
144             repo: context.repo.repo,
145             state: 'open',
146             labels: 'nightly-failure',
147             per_page: 1,
148         });
149         if (existing.data.length > 0) {
150             await github.rest.issues.createComment({
151                 owner: context.repo.owner,
152                 repo: context.repo.repo,
153                 issue_number: existing.data[0].number,
154                 body,
155             });
156         } else {
157             await github.rest.issues.create({
158                 owner: context.repo.owner,
159                 repo: context.repo.repo,
160                 title,
161                 body,
162                 labels: ['nightly-failure'],
163             });
164         }
165     }

```

4. user (2026-06-25T17:13:11.730Z)

```

1     name: CI
2
3     on:
4         push:
5             branches: [master]
6         pull_request:
7
8     permissions:
9         contents: read
10
11     concurrency:
12         group: ci-${{ github.ref }}
13         cancel-in-progress: true
14
15     jobs:
16         # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17         # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
18         # tracked .py, including scripts/ and training/, which the backend-only
19         # py_compile smoke sweep does not reach.
20         lint:
21             name: Ruff lint
22             runs-on: ubuntu-latest
23             steps:
24                 - uses: actions/checkout@v6
25                 - uses: actions/setup-python@v6
26                   with:
27                       python-version: "3.13"
28                 - name: Install ruff
29                   run: python -m pip install ruff
30                 - name: Ruff check
31                   run: ruff check . --output-format github
32
33         # PII / attribution leak guard (tools/leak_scan.py): fails on a committed 32-hex MD5
34         # (the video_id shape) or a personal absolute user path in the shipped-code surface
35         # (backend/training/scripts/tools/eval_baselines; tests/ excluded ? legit fixtures).
Real

```

```

36     # names are scanned only via a gitignored private list locally; CI runs the structural
guard.
37     leak-scan:
38         name: Leak scan (PII / attribution guard)
39         runs-on: ubuntu-latest
40         steps:
41             - uses: actions/checkout@v6
42             - uses: actions/setup-python@v6
43               with:
44                 python-version: "3.13"
45             - name: Leak scan
46               run: python -m tools.leak_scan
47
48     # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
49     # per-module quarantine list ? ratchet by removing entries as modules clean
50     # up. Typed core deps are installed so mypy sees their real signatures
51     # (without them ignore_missing_imports would Any-stub half the checks).
52     typecheck:
53         name: Mypy (lenient baseline)
54         runs-on: ubuntu-latest
55         steps:
56             - uses: actions/checkout@v6
57             - uses: actions/setup-python@v6
58               with:
59                 python-version: "3.13"
60             - name: Install mypy + typed core deps
61               run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
genai numpy
62             - name: Mypy
63               run: mypy backend training
64
65     # Fast guard against the 7fbb0 class of breakage: a SyntaxError or
66     # module-level import error shipping while the suite stays green.
67     # Deliberately does NOT install cv2/torch/numpy ? the smoke tests must hold
68     # on machines without the GPU stack (test_import_smoke stubs what's absent).
69     import-smoke:
70         name: Syntax sweep + import smoke (no GPU stack)
71         runs-on: ubuntu-latest
72         steps:
73             - uses: actions/checkout@v6
74             - uses: actions/setup-python@v6
75               with:
76                 python-version: "3.13"
77             - name: Install minimal deps
78               run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai
79             - name: Run import smoke tests
80               run: python -m pytest tests/test_import_smoke.py -v
81
82     full-suite:
83         name: Full test suite (offline, mocked)
84         runs-on: ubuntu-latest
85         steps:
86             - uses: actions/checkout@v6
87             - uses: actions/setup-python@v6
88               with:
89                 python-version: "3.13"
90                 cache: pip
91             - name: System libs for opencv-python
92               run: |
93                 sudo apt-get update
94                 sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
95             - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
96               # torch/torchvision are NOT in pyproject deps ? their CPU/CUDA wheels

```

```

97         # need an out-of-band index that PEP 621 can't express. Keep installing
98         # them separately here, before the editable install pulls the rest.
99         run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
100         - name: Install package (editable) + dev tooling
101         # Editable PEP 621 install: runtime deps + the `dev` group
102         # (pytest, pytest-cov, httpx for the fastapi TestClient, ruff, mypy).
103         run: python -m pip install -e .[dev]
104         # obsreport (private, soft dep) is absent here: reporting tests skip via
105         # pytest.importorskip ? run_all_tests.py prints which. The coverage floor
106         # is therefore calibrated against THIS lane's number (local runs with
107         # obsreport installed measure a little higher).
108         # Floor calibrated 2026-06-11 against this lane's measured 72% (local
109         # with obsreport: 73%). Ratchet it UP as coverage grows; never down
110         # without an owner decision.
111         - name: Run full suite (with coverage floor)
112         run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
113
114         # Cross-OS determinism guard for the committed golden regression fixtures
115         # (docs/GOLDEN_REGRESSION_FIXTURES.md). The CONTROL replay path is pure-stdlib
116         # (only numpy, no cv2/torch), so this is a FAST, focused job ? it does NOT run
117         # the full suite x3 OSes. It replays the Windows/Linux/macOS runner against the
118         # committed (parse-compared) snapshots, catching any cross-OS float drift beyond
119         # the 1e-6 tolerance before it can flake a contributor's PR on another machine.
120         golden-crossos:
121             name: Golden fixtures cross-OS (${ matrix.os })
122             strategy:
123                 fail-fast: false
124             matrix:
125                 os: [ubuntu-latest, windows-latest, macos-latest]
126             runs-on: ${ matrix.os }
127             steps:
128                 - uses: actions/checkout@v6
129                 - uses: actions/setup-python@v6
130                 with:
131                     python-version: "3.13"
132                 - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
133                 run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
134                 - name: Golden-fixtures characterization (cross-OS determinism guard)
135                 # Also runs the freeze-real / time-origin-reset / M2 refusal-gate tests
(synthetic /
136                 # tmp_path only ? no GPU/network) so the reset metric-invariance (fix #2) and
the
137                 # de-attribution paths get cross-OS coverage and surface any win/mac float
drift.
138                 run: |
139                     python -m backend.eval.golden_fixtures --check
140                     python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
141
142         # Commercial-clean licensing guardrail (MIT/Apache/BSD only ? PACKAGING_PLAN.md P0d).
Scans the
143         # RUNTIME dependency closure (`pip install .` ? no torch/dev) and fails on any
copyleft
144         # (GPL/AGPL/LGPL/SSPL). torch (BSD-3) is installed out-of-band and is not in this
closure. This
145         # automates a guardrail that was human-only; the report is printed for review either
way.
146         license-scan:
147             name: Dependency license scan (no copyleft in runtime deps)
148             runs-on: ubuntu-latest
149             steps:
150                 - uses: actions/checkout@v6

```

```

151     - uses: actions/setup-python@v6
152     with:
153       python-version: "3.13"
154     - name: Install runtime deps + pip-licenses
155       run: |
156         python -m pip install .
157         python -m pip install pip-licenses
158     - name: Report dependency licenses
159       run: python -m piplicenses --format=markdown --with-urls --order=license
160     - name: Fail on any copyleft (GPL/AGPL/LGPL/SSPL) in the runtime closure
161       run: |
162         if python -m piplicenses --format=plain | grep -iE "GPL|SSPL"; then
163           echo "::error::A runtime dependency carries a copyleft (GPL/AGPL/LGPL/SSPL)
license ? the guardrail is MIT/Apache/BSD only. Review the report above."
164         exit 1
165       fi
166       echo "OK: no GPL/AGPL/LGPL/SSPL in the runtime dependency closure."
167
168     # Built-wheel smoke (PACKAGING_PLAN.md P0d): build the wheel, install it NON-editable
from a
169     # non-repo cwd, and launch the console script. Catches packaging regressions the
source suite
170     # can't ? namespace-package resolution (backend has no __init__.py), missing data
files, hidden
171     # imports. (Validated locally: a clean-venv wheel install + CPU torch runs `indexer
--help`.)
172     build-artifact:
173       name: Built wheel smoke (install + launch)
174       runs-on: ubuntu-latest
175       steps:
176         - uses: actions/checkout@v6
177         - uses: actions/setup-python@v6
178         with:
179           python-version: "3.13"
180         - name: System libs for opencv-python
181           run: |
182             sudo apt-get update
183             sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
184         - name: Build the wheel
185           run: |
186             python -m pip install build
187             python -m build --wheel
188         - name: Install the built wheel + CPU torch (out-of-band index)
189           run: |
190             python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
191             python -m pip install dist/*.whl
192         - name: Launch the installed console script from a non-repo dir
193           run: |
194             cd /tmp
195             indexer --help
196             python -c "import backend, training, backend.utils.app_home; print('wheel
import OK')"
197

```

5. assistant (2026-06-25T17:13:25.112Z)

Both files are now read. The finding is verified by the actual file contents.

The finding claims nightly.yml has NO `timeout-minutes` on any job. But the file clearly shows `timeout-minutes` on every single job:

```

- `lint`: line 36 ? `timeout-minutes: 10`
- `typecheck`: line 53 ? `timeout-minutes: 10`
- `full-suite`: line 71 ? `timeout-minutes: 30`

```

- `golden-fixtures`: line 98 ? `timeout-minutes: 15`
- `notify-on-failure`: line 122 ? `timeout-minutes: 5`

These match exactly the values the finding's "Suggestion" field says it "Applied" (lint 10, typecheck 10, full-suite 30, golden-fixtures 15, notify-on-failure 5). So the finding describes its own already-applied fix. The "issue" as stated (no timeouts) is not present in the file being reviewed.

6. user (2026-06-25T17:13:30.795Z)

Structured output provided successfully

7. assistant (2026-06-25T17:13:34.603Z)

The finding is INVALID. The reviewed file nightly.yml already has `timeout-minutes` on all five jobs (lint 10, typecheck 10, full-suite 30, golden-fixtures 15, notify-on-failure 5), which precisely matches the fix the finding claims to have "Applied" ? it describes its own already-shipped change as an open defect.