

claude-windows / 7fce5410-4c57-4ec9-a054-be898bda22da

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\7fce5410-4c57-4ec9-a054-be898bda22da\subagents\workflows\wf_09101709-554\agent-ab9ffd3073f6e0029.jsonl`  
- SHA-256: `ffdcc3fec658a41740a988ad0875d5ac89b5b76bd1a4835fa3ed224ec70b263f`  
- Source modified: `2026-06-25T17:13:15+00:00`  
- Imported at: `2026-07-05T16:48:26+00:00`  
- project: `wf_09101709-554`  
- session_id: `7fce5410-4c57-4ec9-a054-be898bda22da`
```

Transcript

1. user (2026-06-25T17:12:26.126Z)

Adversarially verify this PR-review finding for PR #381 (branch chore/nightly-quality-regression checked out; read .github/workflows/nightly.yml + ci.yml). Be skeptical ? default to INVALID if wrong, already-handled, or a non-issue.

Finding [NIT] Single-purpose, Tier-1-only scope and default-branch caveat are all correctly satisfied ? no change needed
Location: .github/workflows/nightly.yml:1-29, :86-91; PR #381 body
Detail: Positive confirmations, grounded in the files: (a) Single-purpose holds ? git diff is exactly one file, 105 insertions, ci.yml untouched. (b) Tier-1/offline-only holds ? no GPU runner, no real video, no run_regression.py on the corpus; the golden-fixtures lane (92-106) installs only pytest+fastapi+pydantic+numpy and runs backend.eval.golden_fixtures --check (the --check flag exists at golden_fixtures.py:735) plus the two existing test files (tests/test_golden_fixtures.py + tests/test_golden_real_fixtures.py both present). The header (8-12) and the comment at 86-91 state the offline/no-video/no-DB boundary explicitly. (c) Default-branch caveat IS documented in BOTH the workflow header (14-16) AND the PR body ('Scheduled workflows only fire from the repo's default branch ? validate via workflow_dispatch'). (d) name/clarity are good: workflow name 'Nightly quality regression (Tier-1, offline)' and per-job names ('Ruff lint', 'Mypy (lenient baseline)', etc.) match ci.yml and read clearly; concurrency group + permissions: contents: read are set. The workflow_dispatch trigger means the steps CAN be exercised on demand after merge, which is the right mitigation for the no-self-trigger gap. No fix required for this item.
Suggestion: None ? these aspects pass. The only actionable item is the duplication/drift refactor (MEDIUM).

Return isReal + adjustedSeverity (INVALID if not real) + one-paragraph reasoning grounded in the actual files.

2. assistant (2026-06-25T17:12:28.883Z)

I'll verify this finding by examining the actual workflow files and the PR.

3. user (2026-06-25T17:12:30.537Z)

```
chore/nightly-quality-regression  
1a400c0 ci: add Tier-1 nightly quality-regression workflow  
86705c1 chore: ruff format gemini_refine.py + trajectory_hybrid.py (#379)  
b72ac85 fix(B2/P17): surface Gemini provider failures in all callers (#378)
```

4. user (2026-06-25T17:12:31.735Z)

```
.github/workflows/nightly.yml | 105 +++++  
1 file changed, 105 insertions(+)  
---NUMSTAT---  
105      0      .github/workflows/nightly.yml
```

5. assistant (2026-06-25T17:12:35.114Z)

Single-purpose confirmed: exactly one file, 105 insertions, ci.yml untouched. Now let me read the actual workflow file and verify the specific claims.

6. user (2026-06-25T17:12:36.054Z)

```
1     name: Nightly quality regression (Tier-1, offline)
2
3     # Tier-1 nightly quality guard. Per-PR `ci.yml` only fires on push/pull_request;
4     # this runs the SAME offline guards on a schedule so we also catch what per-PR CI
5     # can't: dependency / transitive-dep drift, flakiness over time, and rot during
6     # quiet periods when nothing is pushed.
7     #
8     # Tier-1 == OFFLINE ONLY: no GPU, no real video, no `run_regression.py` on the
9     # real corpus (that's Tier-2 ? owner-gated, needs a self-hosted GPU runner +
10    # the gitignored golden corpus). The lanes here mirror `ci.yml`'s offline jobs
11    # verbatim (same Python pin, same per-lane pip installs) so modules collect/run
12    # instead of silently skipping for a missing dep.
13    #
14    # NOTE: scheduled workflows only fire from the repo's DEFAULT branch, so this
15    # won't actually run on a cadence until it lands on `master`. Use the
16    # `workflow_dispatch` trigger to run it on demand for validation after merge.
17
18    on:
19      schedule:
20        # 03:00 UTC daily.
21        - cron: "0 3 * * *"
22      workflow_dispatch:
23
24    permissions:
25      contents: read
26
27    concurrency:
28      group: nightly-`${{ github.ref }}`
29      cancel-in-progress: true
30
31    jobs:
32      # Lint gate (ruff.toml at repo root) ? mirrors ci.yml::lint.
33      lint:
34        name: Ruff lint
35        runs-on: ubuntu-latest
36        timeout-minutes: 10
37        steps:
38          - uses: actions/checkout@v6
39          - uses: actions/setup-python@v6
40            with:
41              python-version: "3.13"
42          - name: Install ruff
43            run: python -m pip install ruff
44          - name: Ruff check
45            run: ruff check . --output-format github
46
47      # Type gate (mypy.ini at repo root) ? mirrors ci.yml::typecheck. Typed core
48      # deps are installed so mypy sees real signatures (without them
49      # ignore_missing_imports would Any-stub half the checks).
50      typecheck:
51        name: Mypy (lenient baseline)
52        runs-on: ubuntu-latest
53        timeout-minutes: 10
54        steps:
55          - uses: actions/checkout@v6
56          - uses: actions/setup-python@v6
57            with:
58              python-version: "3.13"
```

```

59         - name: Install mypy + typed core deps
60         run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
genai numpy
61         - name: Mypy
62         run: mypy backend training
63
64         # Full offline suite + coverage floor ? mirrors ci.yml::full-suite. CPU-only
65         # torch (no CUDA wheels on a CPU runner), then the editable install pulls the
66         # rest. Floor stays at the ci.yml value (70); never lower it without an owner
67         # decision.
68         full-suite:
69         name: Full test suite (offline, mocked)
70         runs-on: ubuntu-latest
71         timeout-minutes: 30
72         steps:
73         - uses: actions/checkout@v6
74         - uses: actions/setup-python@v6
75           with:
76             python-version: "3.13"
77             cache: pip
78         - name: System libs for opencv-python
79         run: |
80             sudo apt-get update
81             sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
82         - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
83         run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
84         - name: Install package (editable) + dev tooling
85         run: python -m pip install -e .[dev]
86         - name: Run full suite (with coverage floor)
87         run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
88
89         # Offline quality guard: the committed golden regression fixtures
90         # (docs/GOLDEN_REGRESSION_FIXTURES.md) ? synthetic + de-attributed real
91         # trajectory fixtures, NO video / GPU / DB. Mirrors the offline replay in
92         # ci.yml::golden-crossos (single-OS here; ci.yml already covers cross-OS on
93         # every PR ? nightly is about drift-over-time on the default branch). The
94         # CONTROL replay needs only numpy.
95         golden-fixtures:
96         name: Golden fixtures regression (offline quality guard)
97         runs-on: ubuntu-latest
98         timeout-minutes: 15
99         steps:
100        - uses: actions/checkout@v6
101        - uses: actions/setup-python@v6
102          with:
103            python-version: "3.13"
104        - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
105        run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
106        - name: Golden-fixtures characterization (offline)
107        run: |
108            python -m backend.eval.golden_fixtures --check
109            python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
110
111        # Failure visibility ? a silent nightly is useless. Scheduled runs have no PR
112        # to redden and nobody is watching the Actions tab on a quiet day (the exact
113        # "rot during quiet periods" this workflow exists to catch). On ANY guard
114        # failure, open/refresh a single tracking issue so a human is pinged. Uses the
115        # built-in GITHUB_TOKEN (no secrets); `issues: write` is scoped to THIS job
116        # only ? the guard jobs above keep `contents: read`.
117        notify-on-failure:

```

```

118     name: Surface nightly failure (open/refresh tracking issue)
119     needs: [lint, typecheck, full-suite, golden-fixtures]
120     if: failure() && github.event_name == 'schedule'
121     runs-on: ubuntu-latest
122     timeout-minutes: 5
123     permissions:
124       contents: read
125       issues: write
126     steps:
127       - uses: actions/github-script@v7
128         with:
129           script: |
130             const runUrl =
131 `${context.serverUrl}/${context.repo.owner}/${context.repo.repo}/actions/runs/${context.runId}`;
132             const title = 'Nightly quality regression FAILED';
133             const body = [
134               'The Tier-1 nightly offline quality guard failed.',
135               '',
136               ` - Run: ${runUrl}`,
137               ` - Date (UTC): ${new Date().toISOString()}`,
138               '',
139               'One of ruff / mypy / full-suite (cov>=70) / golden-fixtures regressed.',
140               'Open the run for the failing lane. Close this issue once green.',
141             ].join('\n');
142             // Reuse a single open tracking issue instead of spamming one per night.
143             const existing = await github.rest.issues.listForRepo({
144               owner: context.repo.owner,
145               repo: context.repo.repo,
146               state: 'open',
147               labels: 'nightly-failure',
148               per_page: 1,
149             });
150             if (existing.data.length > 0) {
151               await github.rest.issues.createComment({
152                 owner: context.repo.owner,
153                 repo: context.repo.repo,
154                 issue_number: existing.data[0].number,
155                 body,
156               });
157             } else {
158               await github.rest.issues.create({
159                 owner: context.repo.owner,
160                 repo: context.repo.repo,
161                 title,
162                 body,
163                 labels: ['nightly-failure'],
164               });
165             }

```

7. user (2026-06-25T17:12:36.284Z)

```

1     name: CI
2
3     on:
4       push:
5         branches: [master]
6       pull_request:
7
8     permissions:
9       contents: read
10
11     concurrency:
12       group: ci-${{ github.ref }}
13       cancel-in-progress: true

```

```

14
15 jobs:
16   # Lint gate (ruff.toml at repo root): pyflakes/pycodestyle + bugbear.
17   # Also the repo-wide syntax guard ? ruff reports invalid-syntax on every
18   # tracked .py, including scripts/ and training/, which the backend-only
19   # py_compile smoke sweep does not reach.
20   lint:
21     name: Ruff lint
22     runs-on: ubuntu-latest
23     steps:
24       - uses: actions/checkout@v6
25       - uses: actions/setup-python@v6
26         with:
27           python-version: "3.13"
28       - name: Install ruff
29         run: python -m pip install ruff
30       - name: Ruff check
31         run: ruff check . --output-format github
32
33   # PII / attribution leak guard (tools/leak_scan.py): fails on a committed 32-hex MD5
34   # (the video_id shape) or a personal absolute user path in the shipped-code surface
35   # (backend/training/scripts/tools/eval_baselines; tests/ excluded ? legit fixtures).
36   Real
37   guard.
38   # names are scanned only via a gitignored private list locally; CI runs the structural
39   leak-scan:
40     name: Leak scan (PII / attribution guard)
41     runs-on: ubuntu-latest
42     steps:
43       - uses: actions/checkout@v6
44       - uses: actions/setup-python@v6
45         with:
46           python-version: "3.13"
47       - name: Leak scan
48         run: python -m tools.leak_scan
49
50   # Type gate (mypy.ini at repo root): lenient green baseline with an explicit
51   # per-module quarantine list ? ratchet by removing entries as modules clean
52   # up. Typed core deps are installed so mypy sees their real signatures
53   # (without them ignore_missing_imports would Any-stub half the checks).
54   typecheck:
55     name: Mypy (lenient baseline)
56     runs-on: ubuntu-latest
57     steps:
58       - uses: actions/checkout@v6
59       - uses: actions/setup-python@v6
60         with:
61           python-version: "3.13"
62       - name: Install mypy + typed core deps
63         run: python -m pip install mypy pydantic fastapi uvicorn python-dotenv google-
64   genai numpy
65       - name: Mypy
66         run: mypy backend training
67
68   # Fast guard against the 7fbb0 class of breakage: a SyntaxError or
69   # module-level import error shipping while the suite stays green.
70   # Deliberately does NOT install cv2/torch/numpy ? the smoke tests must hold
71   # on machines without the GPU stack (test_import_smoke stubs what's absent).
72   import-smoke:
73     name: Syntax sweep + import smoke (no GPU stack)
74     runs-on: ubuntu-latest
75     steps:
76       - uses: actions/checkout@v6
77       - uses: actions/setup-python@v6
78         with:

```

```

76         python-version: "3.13"
77     - name: Install minimal deps
78       run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai
79     - name: Run import smoke tests
80       run: python -m pytest tests/test_import_smoke.py -v
81
82     full-suite:
83       name: Full test suite (offline, mocked)
84       runs-on: ubuntu-latest
85       steps:
86         - uses: actions/checkout@v6
87         - uses: actions/setup-python@v6
88         with:
89           python-version: "3.13"
90           cache: pip
91         - name: System libs for opencv-python
92           run: |
93             sudo apt-get update
94             sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
95         - name: Install CPU-only torch (no CUDA wheels on a CPU runner)
96           # torch/torchvision are NOT in pyproject deps ? their CPU/CUDA wheels
97           # need an out-of-band index that PEP 621 can't express. Keep installing
98           # them separately here, before the editable install pulls the rest.
99           run: python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
100        - name: Install package (editable) + dev tooling
101          # Editable PEP 621 install: runtime deps + the `dev` group
102          # (pytest, pytest-cov, httpx for the fastapi TestClient, ruff, mypy).
103          run: python -m pip install -e .[dev]
104          # obsreport (private, soft dep) is absent here: reporting tests skip via
105          # pytest.importorskip ? run_all_tests.py prints which. The coverage floor
106          # is therefore calibrated against THIS lane's number (local runs with
107          # obsreport installed measure a little higher).
108          # Floor calibrated 2026-06-11 against this lane's measured 72% (local
109          # with obsreport: 73%). Ratchet it UP as coverage grows; never down
110          # without an owner decision.
111        - name: Run full suite (with coverage floor)
112          run: python run_all_tests.py --cov=backend --cov=training --cov-report=term
--cov-fail-under=70
113
114        # Cross-OS determinism guard for the committed golden regression fixtures
115        # (docs/GOLDEN_REGRESSION_FIXTURES.md). The CONTROL replay path is pure-stdlib
116        # (only numpy, no cv2/torch), so this is a FAST, focused job ? it does NOT run
117        # the full suite x3 OSes. It replays the Windows/Linux/macOS runner against the
118        # committed (parse-compared) snapshots, catching any cross-OS float drift beyond
119        # the 1e-6 tolerance before it can flake a contributor's PR on another machine.
120        golden-crossos:
121          name: Golden fixtures cross-OS (${ matrix.os })
122          strategy:
123            fail-fast: false
124            matrix:
125              os: [ubuntu-latest, windows-latest, macos-latest]
126          runs-on: ${ matrix.os }
127          steps:
128            - uses: actions/checkout@v6
129            - uses: actions/setup-python@v6
130            with:
131              python-version: "3.13"
132            - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
133              run: python -m pip install pytest fastapi uvicorn pydantic python-dotenv google-
genai numpy
134            - name: Golden-fixtures characterization (cross-OS determinism guard)
135              # Also runs the freeze-real / time-origin-reset / M2 refusal-gate tests

```

```

(synthetic /
136         # tmp_path only ? no GPU/network) so the reset metric-invariance (fix #2) and
the
137         # de-attribution paths get cross-OS coverage and surface any win/mac float
drift.
138         run: |
139             python -m backend.eval.golden_fixtures --check
140             python -m pytest tests/test_golden_fixtures.py
tests/test_golden_real_fixtures.py -q
141
142     # Commercial-clean licensing guardrail (MIT/Apache/BSD only ? PACKAGING_PLAN.md P0d).
Scans the
143     # RUNTIME dependency closure (`pip install .` ? no torch/dev) and fails on any
copyleft
144     # (GPL/AGPL/LGPL/SSPL). torch (BSD-3) is installed out-of-band and is not in this
closure. This
145     # automates a guardrail that was human-only; the report is printed for review either
way.
146     license-scan:
147         name: Dependency license scan (no copyleft in runtime deps)
148         runs-on: ubuntu-latest
149         steps:
150             - uses: actions/checkout@v6
151             - uses: actions/setup-python@v6
152               with:
153                 python-version: "3.13"
154             - name: Install runtime deps + pip-licenses
155               run: |
156                 python -m pip install .
157                 python -m pip install pip-licenses
158             - name: Report dependency licenses
159               run: python -m piplicenses --format=markdown --with-urls --order=license
160             - name: Fail on any copyleft (GPL/AGPL/LGPL/SSPL) in the runtime closure
161               run: |
162                 if python -m piplicenses --format=plain | grep -iE "GPL|SSPL"; then
163                     echo "::error::A runtime dependency carries a copyleft (GPL/AGPL/LGPL/SSPL)
license ? the guardrail is MIT/Apache/BSD only. Review the report above."
164                 exit 1
165             fi
166             echo "OK: no GPL/AGPL/LGPL/SSPL in the runtime dependency closure."
167
168     # Built-wheel smoke (PACKAGING_PLAN.md P0d): build the wheel, install it NON-editable
from a
169     # non-repo cwd, and launch the console script. Catches packaging regressions the
source suite
170     # can't ? namespace-package resolution (backend has no __init__.py), missing data
files, hidden
171     # imports. (Validated locally: a clean-venv wheel install + CPU torch runs `indexer
--help`.)
172     build-artifact:
173         name: Built wheel smoke (install + launch)
174         runs-on: ubuntu-latest
175         steps:
176             - uses: actions/checkout@v6
177             - uses: actions/setup-python@v6
178               with:
179                 python-version: "3.13"
180             - name: System libs for opencv-python
181               run: |
182                 sudo apt-get update
183                 sudo apt-get install -y libgl1 libglu1-mesa-dev || sudo apt-get install -y
libgl1 libglu1-mesa-dev
184             - name: Build the wheel
185               run: |
186                 python -m pip install build

```

```

187         python -m build --wheel
188     - name: Install the built wheel + CPU torch (out-of-band index)
189     run: |
190         python -m pip install torch torchvision --index-url
https://download.pytorch.org/whl/cpu
191         python -m pip install dist/*.whl
192     - name: Launch the installed console script from a non-repo dir
193     run: |
194         cd /tmp
195         indexer --help
196         python -c "import backend, training, backend.utils.app_home; print('wheel
import OK')"
```

8. assistant (2026-06-25T17:12:39.140Z)

Now let me verify the specific claims in the finding: the `--check` flag existence, the test files presence, and that the lanes mirror ci.yml.

9. user (2026-06-25T17:12:40.747Z)

```

735:     g.add_argument("--check", action="store_true",
822:     # --check
```

10. user (2026-06-25T17:12:41.612Z)

```

-rw-r--r-- 1 avidu 197609 7293 Jun 25 16:24 tests/test_golden_fixtures.py
-rw-r--r-- 1 avidu 197609 28710 Jun 25 16:24 tests/test_golden_real_fixtures.py
```

11. assistant (2026-06-25T17:12:45.439Z)

Both claims confirmed: `--check` flag exists at golden_fixtures.py:735, and both test files are present. Now let me check the PR body claim about the default-branch caveat.

12. user (2026-06-25T17:12:47.217Z)

```

{"body": "## What\n\nAdds `.github/workflows/nightly.yml` ? a **Tier-1 nightly quality-
regression** workflow that runs the repo's **offline** quality guards on a schedule.\n\n-
**Triggers:** `schedule: cron \"0 3 * * *\"` (03:00 UTC daily) **and** `workflow_dispatch:` (on-
demand).\n\n- **Lanes** (mirror `ci.yml`'s offline jobs verbatim ? same Python 3.13 pin + same
per-lane pip installs so modules actually collect/run instead of skipping for a missing dep):\n
- `ruff check . --output-format github`\n
- `mypy backend training`\n
- `python
run_all_tests.py --cov=backend --cov=training --cov-report=term --cov-fail-under=70`\n
- golden-fixtures regression ? `python -m backend.eval.golden_fixtures --check` + `pytest
tests/test_golden_fixtures.py tests/test_golden_real_fixtures.py -q` (synthetic + de-attributed
real **trajectory** fixtures; **no video / GPU / DB**)\n\n## Why\n\nThe quality/regression
machinery exists but nothing runs on a cadence ? per-PR `ci.yml` only fires on
`push`/`pull_request`. A nightly run of the same offline guards catches what per-PR CI can't:
**dependency / transitive-dep drift, flakiness over time, and rot during quiet periods** when
nothing is pushed.\n\n## Scope ? Tier-1 (offline) only\n\n- **No GPU, no real video, no
`run_regression.py` on the real corpus.** That's **Tier-2** (real-video F1) ? owner-gated, needs
a self-hosted GPU runner + the gitignored golden corpus, and is explicitly out of scope here.\n
Self-contained: **`ci.yml` is untouched**, keeping this a single-purpose change.\n\n## Caveats /
how to validate\n\n- ?? **Scheduled workflows only fire from the repo's default branch.** This
nightly will **not** run on a cadence until it merges to `master`. After merge, validate it on
demand via the **`workflow_dispatch`** trigger (Actions ? `\"Nightly quality regression (Tier-1,
offline)\"` ? Run workflow).\n\n- This PR's own checks are the normal `ci.yml` `pull_request` run
(the new file is just YAML; it doesn't self-trigger on the PR).\n\n## Follow-ups (NOT in this
PR)\n\n- **M4 quality-floor** ? add a per-slug F1 floor in
`eval_baselines/regression_baseline.json` + a windowing fingerprint (CONTROL-only) so the
golden-fixtures regression asserts an actual **F1 floor** rather than just exercising the flow.
See `docs/NEXT_STEPS.md` / `docs/GOLDEN_REGRESSION_FIXTURES.md`.\n\n- **Tier-2** ? real-video F1
via a self-hosted GPU runner is a separate, **owner-gated** effort.\n\n? Generated with [Claude
Code](https://claude.com/claude-code)\n", "files": [{"path": ".github/workflows/nightly.yml", "addit
```



```
ions":105,"deletions":0,"changeType":"ADDED"]},"title":"ci: add Tier-1 nightly quality-  
regression workflow"}
```

13. assistant (2026-06-25T17:12:51.502Z)

All claims in the finding are verified against the actual files. Let me cross-check the lane-naming and offline-only claims one more time, and confirm there's no GPU/real-video/run_regression.py reference in the nightly workflow.

14. user (2026-06-25T17:12:52.915Z)

```
8:# Tier-1 == OFFLINE ONLY: no GPU, no real video, no `run_regression.py` on the  
9:# real corpus (that's Tier-2 ? owner-gated, needs a self-hosted GPU runner +  
10:# the gitignored golden corpus). The lanes here mirror `ci.yml`'s offline jobs  
65: # torch (no CUDA wheels on a CPU runner), then the editable install pulls the  
82:     - name: Install CPU-only torch (no CUDA wheels on a CPU runner)  
91: # trajectory fixtures, NO video / GPU / DB. Mirrors the offline replay in  
104:     - name: Install minimal deps (no GPU stack ? the CONTROL replay needs only numpy)
```

15. assistant (2026-06-25T17:13:03.450Z)

The only references to GPU/real-video/run_regression are in comments documenting the *exclusion* of those (Tier-2 boundary), and CPU-only torch installs. No actual GPU runner, real-video step, or `run\_regression.py` invocation exists. Every positive confirmation in the finding holds.

Let me verify the finding's own internal consistency on the line references it cites (it cites lines :1-29, :86-91 for nightly.yml, but the actual offline comment is at :8-12 and the golden-fixtures comment is at :89-94 ? minor line-number drift in the finding's citation, but the substance is correct).

The finding is a NIT-level "no change needed / pass confirmation". All four sub-claims (a?d) check out:

- ***(a) Single-purpose***: confirmed ? `git diff --numstat` shows exactly `105 0 .github/workflows/nightly.yml`, ci.yml untouched.
- ***(b) Tier-1/offline-only***: confirmed ? no GPU runner, no real video step, no `run\_regression.py` invocation; golden-fixtures lane installs only `pytest+fastapi+uvicorn+pydantic+python-dotenv+google-genai+numpy` and runs ``backend.eval.golden_fixtures --check`` (flag exists at `golden_fixtures.py:735`) + the two present test files. The ``--check``-only minor note: the finding said it installs "only `pytest+fastapi+pydantic+numpy`" ? the actual install list is slightly broader (also `uvicorn`, `python-dotenv`, `google-genai`), but this is an understatement that doesn't change the offline conclusion.
- ***(c) Default-branch caveat***: documented in both the workflow header (lines 14-16) and the PR body.
- ***(d) name/clarity/permissions/concurrency***: confirmed ? workflow name, per-job names match ci.yml, ``permissions: contents: read``, concurrency group all present.

16. user (2026-06-25T17:13:11.282Z)

Structured output provided successfully

17. assistant (2026-06-25T17:13:15.733Z)

Verified. The finding is a real NIT-level positive confirmation, and all four sub-claims (a?d) hold against the actual files:

- ``github/workflows/nightly.yml`` ? the only changed file (105 insertions, 0 deletions); ``ci.yml`` untouched.
- The golden-fixtures lane runs ``backend.eval.golden_fixtures --check`` (flag confirmed at ``backend/eval/golden_fixtures.py:735``) plus the two present test files.
- Default-branch caveat documented in both the workflow header and the PR body.

Severity: NIT, isReal: true. Two trivial wording slips in the finding (line-number citations off

by a few lines, and the golden lane's pip list under-stated) don't change any conclusion.