

claude-windows / accd47af-95a3-46a7-9b51-b3eb5066a777

Metadata

- Source: `claude-windows`
- Kind: `claude`
- Source file: `C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\accd47af-95a3-46a7-9b51-b3eb5066a777\subagents\workflows\wf\_793a5d86-1c4\agent-a6030653e9c572460.jsonl`
- SHA-256: `b2a4ff95b28a7d2f8006754e950070c8d0ee4f0a40d7c68ca011fd70277f17fe`
- Source modified: `2026-07-02T08:20:47+00:00`
- Imported at: `2026-07-05T16:48:24+00:00`
- project: `wf\_793a5d86-1c4`
- session_id: `accd47af-95a3-46a7-9b51-b3eb5066a777`

Transcript

1. user (2026-07-02T08:19:35.490Z)

You are an ADVERSARIAL verifier. Independently re-check the fate verdict for OPEN PR #420 ("docs: reconcile stale trackers/blurbs (R1-13/R1-14/R1-01-doc/DATA_IN_GCS)", branch origin/docs/reconcile-trackers-2026-07-02) on Khelsutra/badminton-highlight-indexer. Try to REFUTE the analyst.

Analyst's verdict: recommended_fate=TAKE_OVER_REBASE; superseded_by=421 (and 423) (confidence high).

Analyst's master_evidence: PR branched at e410136, BEFORE #421 (d80aa02) and #423 (e517475), which reconciled the SAME audit findings. SUPERSEDED hunks (4 files master changed since base ? `git diff --numstat e410136 HEAD`: BACKLOG 4+/4-, DOC_STATUS 5+/3-, README 4+/3-, NEXT_STEPS 2+/1-): (1) docs/BACKLOG.md:34 already "? DONE 2026-06-27 (#386)... supervision cap lifted (supervision>=0.21.0)" and lines 25-26 already reconcile the app-factory (create_app #370, no lifespan) + WSL/Gemini-GC (#380 shipped, retention remains) rows; #420's "? PARTIAL" versions restate the same facts. (2) docs/DOC_STATUS.md:47 already has the DECISION_SNAPSHOT_REGRESSION row (#420 claimed it was "missing entirely"), :43 CODE_CLEANUP row reconciled, :26 COMPUTE_DECOUPLED "M1?M9 LARGELY SHIPPED", :3 "Last full review: 2026-07-02". (3) docs/README.md:42 already off "DRAFT, owner-gated"; :44 already off "decisions in \$2 need locking". (4) docs/NEXT_STEPS.md:1,5,6 already rewritten to "M0?M9 LARGELY SHIPPED"; stale "owner M0 design sign-off pending" gone. STILL-UNAPPLIED / NOT superseded (3 files master never touched, numstat=UNCHANGED): (a) backend/pipeline/segmenters/yolo_hybrid.py:12-13 STILL "supervision's MIT-licensed ByteTrack" ? #386 never fixed the docstring. (b) docs/DATA_IN_GCS.md:120-124 \$5 gap table STILL ? for trajectories/manifest/labels, contradicting \$3:68 ("Phase-1 landed 2026-06-22"), \$7:188-192, \$8:199. (c) docs/CODE_CLEANUP_AUDIT_2026-06-24.md:10 STILL "\$7 progress tracker is the source of truth" though the tracker header is "## 6. Deliverables & progress tracker ? source of truth" (:174); P10 row :198 still "~1450 lines / ? / ?" with no #390 note. #420's README fact-fixes ("35 findings" and "\$6") are correct & unapplied (audit \$9:261 "total of 35"; tracker is \$6). Analyst's conflict_nature: CONFLICTING on exactly the 4 files master edited since the PR base (BACKLOG.md, DOC_STATUS.md, README.md, NEXT_STEPS.md) ? and on those files the conflict is because master ALREADY performed the same reconciliation via #421/#423 (superseded), not adjacent edits. The other 3 touched files (yolo_hybrid.py, DATA_IN_GCS.md, CODE_CLEANUP_AUDIT_2026-06-24.md) are UNCHANGED on master since the base, so those hunks would apply cleanly and carry real, still-unapplied value.

Analyst's rationale: Roughly 4/7 files are fully superseded by #421/#423 (same audit findings R1-13/R1-14, same reconciliations) and must be dropped ? the DOC_STATUS DECISION_SNAPSHOT row #420 calls "missing" already exists, ByteTrack is already flipped, NEXT_STEPS "Where we are" and README blurbs are already reconciled. But 3 files carry genuine, still-unapplied fixes that master does NOT yet have: (1) the yolo_hybrid.py docstring still wrongly says "supervision's MIT-licensed ByteTrack" (a licensing-accounting error #386 left behind); (2) DATA_IN_GCS \$5's gap table still self-contradicts \$3/\$7/\$8 by showing metadata as ? though it landed in-bucket 2026-06-22; (3) CODE_CLEANUP_AUDIT still has the \$7?/\$6 anchor bug and a stale P10 row (no #390-stale note, wrong main.py size). These three are worth keeping. Not CLOSE_SUPERSEDED (real residual value) and not TAKE_OVER_CLEAN (4 files conflict + a stale line-count). Recommend re-cutting a small PR from origin/master that keeps only the yolo_hybrid.py, DATA_IN_GCS.md, and

CODE_CLEANUP_AUDIT hunks, drops the four superseded files, and refreshes the main.py size from 2198 to the current 2360.

Already MERGED onto master (verify against actual master code, do not assume):

- #422 = fix for --output-dir clobbering config.json output.output_dir (CODE_MAP updated for it)
- #423 = docs: health-remediation plan added + DOC_STATUS/README/NEXT_STEPS registered; audit marked snapshot
- #424 = P2: /api/process idempotent by content hash (#340 re-bill fix)
- #425 = P3: /api/source resolves remote source into output-scoped source_cache, reuses across Range requests, warms annotation proxy
- #426 = P4: bounded ffmpeg/ffmpeg subprocess timeouts on serving paths
- #427 = P5: clean remote-input scratch dirs after process/compile/assets/CLI; durable reports
- #428 = P1/P6: backend/storage/gcs.py lazy-import mypy typing fix

master HEAD = d00b4db "fix gcs storage lazy import typing (#428)"

Working dir: C:/Users/avidu/Projects/khelsutra-guru/badminton-highlight-indexer on master.
SHARED checkout ? READ-ONLY, no branch checkout, no working-tree changes.

Your job ? verify by looking at code yourself, do NOT take the analyst's word:

- If the verdict is CLOSE_SUPERSEDED: the risk is closing a PR whose work is NOT actually on master (losing real work). Open the specific master file:line the analyst cited AND the PR diff (gh pr diff 420 -R Khelsutra/badminton-highlight-indexer). Confirm master truly implements the SAME intent (not just a similarly-named symbol). If the PR does anything master lacks (extra tests, a stricter guard, a different file), flag it ? that means NOT fully superseded.
- If the verdict is TAKE_OVER_*: the risk is that the work IS already on master (would be redundant/conflicting) OR the change is buggy. Confirm master still lacks it, and sanity-check the change is correct.
- Default to skepticism: only agree=true if the evidence in code clearly supports it.

Return ONLY the structured object.

2. assistant (2026-07-02T08:19:37.760Z)

I'll independently verify the analyst's claims by examining master code directly.

3. user (2026-07-02T08:19:41.319Z)

d00b4db fix gcs storage lazy import typing (#428)
---base merge-base---
e4101365cb5f3a608e76f8baf1d2e0d65ce1738e

4. user (2026-07-02T08:19:43.100Z)

<persisted-output>

Output too large (33.5KB). Full output saved to: C:\Users\avidu\claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\accd47af-95a3-46a7-9b51-b3eb5066a777\tool-results\bf2zt3xmd.txt

Preview (first 2KB):

```
diff --git a/backend/pipeline/segmenters/yolo_hybrid.py
b/backend/pipeline/segmenters/yolo_hybrid.py
index a41c83b..556a653 100644
--- a/backend/pipeline/segmenters/yolo_hybrid.py
+++ b/backend/pipeline/segmenters/yolo_hybrid.py
@@ -9,8 +9,10 @@
```

Stage 1 originally used the AGPL-licensed ultralytics YOLO. To keep the hosted service free of copyleft obligations (see docs/MONETIZATION_AUDIT.md and ADR-005), detection uses torchvision's permissively-licensed (BSD) detectors and the -association/tracking that YOLO.track() used to bundle in is provided by -supervision's MIT-licensed ByteTrack.
+association/tracking that YOLO.track() used to bundle in is provided by ByteTrack
+via the standalone Apache-2.0 ``trackers`` package (D13/P12 migration, #386 ?
+previously supervision's ``sv.ByteTrack``, removed upstream in supervision 0.30;

+see ``person\_detector.make\_tracker`` for the behaviour-pinning details).

The registry key stays "yolo_hybrid" for back-compat with existing configs and the DB `detector` tags, even though no YOLO is involved anymore.

diff --git a/docs/BACKLOG.md b/docs/BACKLOG.md

index 29ecd98..b02d4c7 100644

--- a/docs/BACKLOG.md

+++ b/docs/BACKLOG.md

@@ -5,7 +5,7 @@ Single tracker for work we consciously deferred, so nothing gets lost between se
> Maintained as a version-controlled doc for now. Say the word and I'll mirror any/all of these
to **GitHub issues** instead (or in addition).

> Status legend: ? blocks something · ? do when triggered · ? nice-to-have/cleanup.

-_Last updated: 2026-06-23._

+_Last updated: 2026-07-02 (reconciled the three 2026-06-10-sweep rows below against shipped
code: #370/#380/#386)._

@@ -19,8 +19,8 @@ _Last updated: 2026-06-23._

- ? **Define the ship-gate F1@tIoU target** ? "very high P/R" must be operationalized (floor =
beat heuristic LOVO 0.480; learned is now 0.626 post metric-fix). **Revisit:** at M0.4 harness.

Engineering hardening deferred from the 2026-06-10 audit sweep

...

</persisted-output>

5. assistant (2026-07-02T08:19:46.618Z)

The first PR hunk (yolo_hybrid.py) is visible. Let me verify each claim against master. Let me
check the numstat and the specific "still-unapplied" files first.

6. user (2026-07-02T08:19:52.053Z)

=== numstat base->HEAD for the 7 files ===

4	4	docs/BACKLOG.md
5	3	docs/DOC_STATUS.md
2	1	docs/NEXT_STEPS.md
4	3	docs/README.md

7. user (2026-07-02T08:19:52.985Z)

=== yolo_hybrid docstring on master ===

(`PersonDetector`, a reusable rally-cue producer ? fusion P1) so the SAME person cue
can be consumed by the future fusion segmenter; this segmenter just orchestrates
chunking, the AI-provider handoff, and DB population around it. Behaviour is unchanged.

Stage 1 originally used the AGPL-licensed ultralytics YOLO. To keep the hosted
service free of copyleft obligations (see docs/MONETIZATION_AUDIT.md and ADR-005),
detection uses torchvision's permissively-licensed (BSD) detectors and the
association/tracking that YOLO.track() used to bundle in is provided by
supervision's MIT-licensed ByteTrack.

The registry key stays "yolo_hybrid" for back-compat with existing configs and the
DB `detector` tags, even though no YOLO is involved anymore.

"""

```
import concurrent.futures
import logging
```

8. assistant (2026-07-02T08:19:57.656Z)

Confirmed: numstat shows exactly 4 files changed since base (BACKLOG, DOC_STATUS, NEXT_STEPS,
README) ? the other 3 (yolo_hybrid.py, DATA_IN_GCS.md, CODE_CLEANUP_AUDIT) are UNCHANGED on

master. And the yolo_hybrid.py docstring on master STILL says "supervision's MIT-licensed ByteTrack" ? confirming the analyst's key claim (a).

Let me verify claims (b) DATA_IN_GCS and (c) CODE_CLEANUP_AUDIT, plus spot-check the 4 "superseded" files.

9. user (2026-07-02T08:20:00.237Z)

```
=== DATA_IN_GCS $5 gap table (110-130) ===
A bare/relative path resolves under a configured `input_root`/`input_backend`
(`backend/storage/ref.py::resolve_input_uri`), so workflows can be written bucket-relative.

---
```

5. The gap ? what's still local-only (why the box is still a blocker)

All of this is **tiny** except the originals ? the metadata that blocks cloud eval is < 25 MB total:

Data	Where it is now	Size	In bucket?	Needed for
20 WASB trajectory CSVs	`C:\?\Annotation Setup\Collect\Trajectories\`	17.5 MB	? none	eval/litmus (cached). *(No `cached` train source exists ? train synths or detector-extracts.)*
golden manifest	`output/human_lovo_manifest.json` (C:-absolute paths)	4.5 KB	? eval	corpus definition
original-6 labels	`F:\[Khelsutra]\Training\Golden Labelled\*.rallies.csv`	~12 KB	?	train + eval
7 new-clip labels	repo `output/*.rallies.csv` (gitignored) *** F: archive (backed up 2026-06-22)	~12 KB	? partial/inconsistent	train + eval
source videos / proxies	F: archive + Google Drive golden folder	100s of GB	? 3 only	infer + train (real WASB extraction)

```
> The corpus *metadata* (trajectories + labels + manifest, ~25 MB) is what actually blocks GPU-free
> cloud eval and the served Gate-A litmus ? and it's trivially small to migrate. The **videos** are the
> only large migration, and they already exist on **Google Drive** (the golden-corpus folder), so they
> can go **Drive ? bucket from a cloud box**, keeping the local machine out of the path entirely.
=== $3 ~68 ===
---
```

3. What's actually in the bucket today (2026-06-22)

```
> **Phase-1 landed 2026-06-22:** added `trajectories/` (15), 15 canonical
`labels/<date>_<name>.rallies.csv`
> (= original-6 + `GX010137`/`GX010141` + 7 new; the 7 old `_proxy`/bare dupes were deleted), and
> `manifests/golden.json`. The table below is the **pre-Phase-1** snapshot; the remaining gap is
**videos**
> (14 of 15) ? see $6/$7b/$8.
```

10. user (2026-07-02T08:20:02.795Z)

```
=== CODE_CLEANUP_AUDIT line 10, 174, 198, 261 ===
> **Status:** `PHASE 1 COMPLETE` ? 14 items shipped + D2 app-factory + test fixes . **Owner:** Avi .
**Created:** 2026-06-24 . **Last updated:** 2026-06-25 (DoD/line-count accuracy fixes ? see Changelog)
> **Lifecycle:** `DRAFT ? IN PROGRESS ? DONE ? archived` (move to `docs/archives/past_projects/` when DONE)
> **Tracking anchors:** $7 progress tracker is the source of truth; indexed in `docs/README.md`.
> **Relation to existing docs:** Sibling of `docs/CODE_AUDIT_AND_TEST_HARDENING.md` (the
```

2026-06-10 hardening sweep).

>
---174---

6. Deliverables & progress tracker ? **source of truth**

Legend: ? Todo · ? In progress · ? Done · ? Blocked/gated. **One small PR per row**, branched from `origin/master`, no stacking.

---198 P10---
| **P8** | **B2: Error-flagging in Gemini provider metrics** (Phase 2) | P7 | No | ? |
[#373](https://github.com/Khelsutra/badminton-highlight-indexer/pull/373) |
| **P9** | **B1: Remove fabricated data from motion segmenter** ? random
server/receiver/shot_count/speed ? `None`; events dropped | P7 | No | ? |
[#385](https://github.com/Khelsutra/badminton-highlight-indexer/pull/385) |
| **P10** | **D1: Split `main.py` god-module** (~1450 lines ? smaller modules) | P7 | No | ? | ?
|
| **P11** | **O1: Gemini remote-file GC** ? per-call `\_safe\_delete` (#373) + startup orphan
sweep scoped to our `UPLOAD\_TAG`-tagged files, grace-gated, HTTP-timeout-bounded | P7 | No | ? |
[#380](https://github.com/Khelsutra/badminton-highlight-indexer/pull/380) |
| **P12** | **D13: ByteTrack deprecation** ? migrated `sv.ByteTrack` ?
`trackers.ByteTrackTracker` (git dep; PyPI wheel broken) | P7 | No | ? |
[#386](https://github.com/Khelsutra/badminton-highlight-indexer/pull/386) |
---261---

9. References

- **Audit source:** 33 findings from a deep whole-repo read on 2026-06-24, plus 2 absorbed from `docs/BACKLOG.md` (D13, D14) for a total of 35. Two additional BACKLOG items (D2, O1) were already independently found and are cross-referenced in §3e.
- **Prior hardening sweep:** `docs/CODE\_AUDIT\_AND\_TEST\_HARDENING.md` (completed 2026-06-14) ? the previous audit-driven hardening round. This project is the next wave.
- **BACKLOG.md** `docs/BACKLOG.md` ? deferred items from the 2026-06-10 sweep; §3e maps the absorbed items.

11. assistant (2026-07-02T08:20:10.343Z)

All three "still-unapplied" claims are confirmed on master:

(a) yolo_hybrid.py:12-13 STILL "supervision's MIT-licensed ByteTrack" ? (confirmed above)

(b) DATA_IN_GCS §5 gap table STILL shows ? for trajectories/manifest and ? for labels, while §3:66 says "Phase-1 landed 2026-06-22" adding those to the bucket ? self-contradiction confirmed ?

(c) CODE_CLEANUP_AUDIT:10 STILL says "§7 progress tracker is the source of truth" but the tracker header is "## 6. Deliverables & progress tracker" (line 174) ? anchor bug confirmed ?. P10 row (line 198) still "~1450 lines ? smaller modules ... ? ... ?" with no #390 note ?. And §9:261 confirms "total of 35" (README fact-fix correct) ?.

Now let me verify the "superseded" claims on the 4 changed files, and check the main.py line count claim (2198 vs 2360).

12. user (2026-07-02T08:20:15.836Z)

=== main.py current line count ===
2360 backend/main.py

13. user (2026-07-02T08:20:17.151Z)

=== DOC_STATUS.md master: DECISION_SNAPSHOT, line 3, 26, 43, 47 ===
3:> **Status:** ? **LIVING** ? the single board for doc **lifecycle + archival due-diligence**.
Last full review: 2026-07-02.
33:| `COMPUTE\_DECOUPLED\_SERVING/` | **IN PROGRESS** ? M0 approved + **M1?M9 LARGELY SHIPPED**

(#279?#290; M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF); Gen-0 weights pushed (`weights/0.1.0-l4/`) | No | M11 flywheel + M12 gdrive-api (owner residual: counsel/prices/OAuth/real-run within the D17 cap) |
47:| `CODE\_CLEANUP\_AUDIT\_2026-06-24.md` | ****IN PROGRESS**** ? Phase 1 COMPLETE (14 items shipped); Phase 2 nearly done (P8/P9/P11-P15/P17 merged #373-#386). Remaining: P10 (main.py split, stalled ? PR #390 CONFLICTING) + P16+ (Phase 3). | No | Re-cut P10 split; finish P16+ Phase 3 items |
49:| `DECISION\_SNAPSHOT\_REGRESSION.md` | ****DRAFT ? decisions LOCKED**** (D1?D11 locked 2026-06-25, #383). Zero code shipped for P1?P7. | No | Start P1+P2 (schema + stitch-plan factor); de-risk the P10 main.py split |
60:| `PLATFORM\_ARCHITECTURE.md` | Strategy/research | The `ComputeBackend`/`StorageBackend` registries ? realized by COMPUTE_DECOUPLED_SERVING |
84:****ADR-001 ? ADR-014**** (14, all ratified). The compute/serving lineage: ****ADR-005 ? 008 ? 009 ? 010 ? 011 ? 012 ? 013 ? 014****. ADR-014 was ****RATIFIED 2026-06-21**** (owner M0 sign-off for COMPUTE_DECOUPLED_SERVING).
94:6. ****? RESOLVED (2026-06-21) ? `COMPUTE\_DECOUPLED\_SERVING` M0 approved + M1?M4 SHIPPED**** (#279/#280/#282/#283); all \$8 gating Qs locked (D7?D17); M5+ default-OFF code batch authorized (owner residual: counsel ToS/DPA, real GCP prices, OAuth app, real cloud run within the D17 \$100/?10k cap).

14. assistant (2026-07-02T08:20:21.994Z)

Confirmed on master DOC_STATUS.md:

- Line 3: "Last full review: 2026-07-02" ? (matches analyst, means #421/#423 already reconciled)
- Line 33: COMPUTE_DECOUPLED "M1?M9 LARGELY SHIPPED" ?
- Line 47: CODE_CLEANUP row reconciled with P10/P390 ?
- Line 49: DECISION_SNAPSHOT_REGRESSION row EXISTS ? ? #420 claimed it was "missing entirely" but master already has it. Superseded confirmed.
- main.py is 2360 lines (not 2198) ? ? the analyst's stale-line-count observation is correct.

Let me verify the README and BACKLOG "superseded" claims to be thorough, and confirm the #421/#423 commits reconciled these.

15. user (2026-07-02T08:20:24.129Z)

=== README lines 40-46 ===

- CODE_AUDIT_AND_TEST_HARDENING.md ? ****COMPLETED**** (2026-06-14; T1?T5 + full loop + verification per #162). Living record of the June 2026 plumbing / contract / architectural audit, 626-test infrastructure, positive test plans for the 5 identified issues, multi-layer regression protocol (including ablation/experiment layer), and pre/post-change verification checklist. See `docs/archives/CODE\_AUDIT\_AND\_TEST\_HARDENING-COMPLETED-2026-06-14.md` and `docs/archives/checkpoints/2026-06-14-hardening-loop-complete.md` for the final archived state. (Top-level is now pointer stub per review; full content in archives/.)
- CODE_CLEANUP_AUDIT_2026-06-24.md ? ****? Tracked project** (Phase 1 COMPLETE, Phase 2 nearly done):****** the next-wave cleanup project ? 33 audit findings across bugs, tech debt, operational risks, and quick wins. Phase 1 = pure refactorings (typed-config migration, app-factory, segmenter validator fix); Phase 2+ = logic fixes. \$7 progress tracker is the source of truth; sibling of the completed hardening sweep above.
- KHELSUTRA_GURU_HEALTH_REMEDIATION_PLAN.md ? ****? Tracked project** (IN PROGRESS, 2026-07-02):****** live execution board for the five-repo `kheksutra-guru` audit cleanup. Extends the 2026-07-02 audit report, records post-sync completed items and pushback, then sequences the work into small PR waves with coverage and UI-screenshot requirements. \$7 progress tracker is the source of truth.
- DECISION_SNAPSHOT_REGRESSION.md ? ****? Tracked project** (decisions D1?D11 LOCKED 2026-06-25, #383; implementation P1?P7 pending):****** a deterministic ****quality-neutral proof**** ? freeze + compare the pipeline's ****decision output**** (rally segments + stitch/concat plan) off a cached trajectory, GPU-free, in CI. The behavioral proof that lets a ***genuine*** (non-AST-identical) refactor be confidently called quality-neutral. ****Never** hashes the encoded reel****** (non-deterministic). Extends the existing characterization snapshot (`test\_golden\_fixtures.py`) from segmentation to the stitch plan; peer of the M4 quality-floor + GPU `nightly\_reelcount.py`. \$7 tracker is the source of truth.
- ****Field-ops / customer-feedback**** (recruiting flyer · field-visit playbook · PMF) ?
[`archives/field-pmf/`](archives/field-pmf/) (archived 2026-06-21; PMF validation still an active ***business*** thread ? see `NEXT\_STEPS.md`).

```
- **Roora annotation-vendor** engagement (labeling brief · vendor bake-off) ? [`archives/roora-vendor/`](archives/roora-vendor/) (archived 2026-06-21).
=== BACKLOG 24-40 ===
- ? **shlex-quote WSL config values** ? **DONE 2026-06-11 (#97):**
`detectors/base.py::wsl_tilde_quote` quotes everything after a leading `~/`, so quoting and
tilde expansion coexist; applied to every `~`-capable arg in both runners;
`tests/test_wsl_tilde_quoting.py` pins the command shapes (no quoted tilde can ever reach
`_wsl_bash`). Found live: the naive quoting had silently broken wasb_hybrid staging for the
default `~/clips`. Residual for tenancy-era: quote the remaining trusted-config interpolations
(repo_dir/conda_env) if config ever becomes per-tenant.

---
```

Monetization / licensing

(source: `archives/MONETIZATION\_AUDIT.md` §7, ADR-005)

```
- ? **Ultralytics Enterprise pricing** ? get a quote **only if** the torchvision/ByteTrack
detector (PR #8) proves insufficient on real footage. Currently *not needed* (swap done). Quote-
only pricing. **Revisit:** if detector quality forces reconsidering YOLOv8.
- ? **Codec patent thresholds (AVC/HEVC/AAC)** ? confirm exact Via-LA (AVC) / Access Advance
(HEVC) royalty rates & free thresholds **with legal counsel** before relying on the "small SaaS
is under threshold" reading. Mitigation already chosen: emit royalty-free AV1/Opus output.
**Revisit:** before commercial launch.
- ? **Per-checkpoint weight/dataset licenses** ? if we borrow pretrained weights (YOLOX,
ActionFormer, WASB, TrackNetV3), the *code* license ? the distributed *weights/dataset* terms.
Either train our own or confirm each checkpoint. **Revisit:** when adopting any borrowed
weights.
- ? **Real GPU $/video & latency** ? benchmark the three LLM-step options (offline ActionFormer
vs self-hosted Qwen2.5-VL-7B vs Gemini paid-tier) on the target VM before pricing the product.
**Revisit:** when choosing the production LLM step.
```

Detector swap follow-ups

(source: PR #8, ADR-005)

```
- ? **`supervision` ByteTrack deprecation** ? **DONE 2026-06-27 (#386):** migrated to
`ByteTrackTracker` from the separate `trackers` package
(`backend/pipeline/detectors/person_detector.py`); the `<0.30.0` supervision cap is lifted
(`supervision>=0.21.0`); `trackers @ git+...@2.5.0` pinned in `requirements.txt`.
- ? **Real-video validation of the swap** ? the torchvision + ByteTrack Stage-1 is unit-tested
with mocks but **not yet validated on actual footage** for detection/tracking quality. (ADR-004
separately found player pixel-motion is a weak rally proxy on broadcast video, so expect
tuning.) **Revisit:** when a real test clip + the pipeline are run together.
```

16. user (2026-07-02T08:20:27.513Z)

=== NEXT_STEPS 1-8 ===

Next steps (live) ? refreshed 2026-06-21

> **How to read:** this top section is the **current cross-track state + prioritized pending list** (refreshed 2026-06-21, replacing the stale 2026-06-13 header). The detailed **owned-model & rally-quality track** plan (2026-06-13) is **preserved below, still valid** ? it is now *one of several* active tracks. Live blow-by-blow: memory `current-state.md`. Doc lifecycle/hygiene: `docs/DOC\_STATUS.md`.

? Where we are (2026-06-21)

DECOUPLED_COMPUTE delivered on GCP + archived (#270); **compute = GCP only** (ADR-013, DO dropped). **`COMPUTE\_DECOUPLED\_SERVING`** subproject spun up (#272, **ADR-014**) ? one pure core + local/cloud profiles; **M0?M9 LARGELY SHIPPED** (ADR-014 ratified 2026-06-21; M1?M4 #279?#283, M5 #286, M6 #287, M8 #288, M9-auth #290 all merged, default-OFF). **Per-video workspace P0** shipped (#264); **doc-status register + archival due-diligence** in place (#273). The **khelsutra site + per-rally UI** ship in parallel (sibling track). The **rally-quality / owned-model** track (below) is the core ML work, gated on **growing the golden corpus**.

```
> **? LAUNCH DECISION ? NO-GO (owner, 2026-06-30).** The rally-detection product is **not**
cleared for an external/commercial launch; we stay in **development**. This is the expected
default for the dev phase. Two unblockers gate a future GO: **(1)** set + meet the **ship-gate
target (#172)** ? currently unset, needs corpus growth (and the R2 tolF1 ablation, item 8);
**(2)** resolve the **WASB shuttle-detector weight provenance** (commercial-clean licensing).
P2b (#396) was a do-no-harm quality increment, **not** a launch gate ? merging dev PRs ?
launching. Re-decide only when both unblockers clear.
=== commits touching these 4 docs since base ===
e517475 docs: track khelsutra-guru health remediation plan
d80aa02 docs: truth-sync ? audit report + fix stale state in CLAUDE.md, BACKLOG, DOC_STATUS,
README, NEXT_STEPS (#421)
```

17. assistant (2026-07-02T08:20:36.754Z)

All claims verified. The 4 changed docs were reconciled by #421 (d80aa02) and #423 (e517475), exactly as the analyst said:

- README:44 ? DECISION_SNAPSHOT row exists, no "DRAFT owner-gated / decisions need locking" stale text (now "decisions D1?D11 LOCKED") ? supersedes #420's fixes
- BACKLOG:34 ? "? DONE 2026-06-27 (#386) ... supervision cap lifted (supervision>=0.21.0)" ? supersedes #420
- NEXT_STEPS:6 ? "M0?M9 LARGELY SHIPPED", no stale "owner M0 design sign-off pending" ?
- DOC_STATUS ? DECISION_SNAPSHOT row present, "Last full review 2026-07-02" ?

Note one nuance: README:42 (the CODE_CLEANUP line) on master still says "33 audit findings" and "\$7 progress tracker is the source of truth" ? #420's README fact-fix to "35 findings"/"\$6" is genuinely unapplied there too. But #420's README diff also touches the superseded blurbs, so the README file overall conflicts. The analyst's core split (4 superseded / 3 clean-carry) holds, and the residual value in the 3 unchanged files is real and confirmed.

The verdict TAKE_OVER_REBASE is well-supported: real residual value exists (rules out CLOSE_SUPERSEDED), and 4 files conflict via same-intent reconciliation plus a stale line count (rules out TAKE_OVER_CLEAN). I agree.

18. user (2026-07-02T08:20:47.829Z)

Structured output provided successfully