

claude-windows / c1fcb4aa-537d-41fd-858e-53f93349bf5a

Metadata

```
- Source: `claude-windows`  
- Kind: `claude`  
- Source file: `C:\Users\avidu\.claude\projects\C--Users-avidu-Projects-badminton-highlight-indexer\c1fcb4aa-537d-41fd-858e-53f93349bf5a\subagents\workflows\wf_7049e1ed-b93\agent-a0f976262d0bcf4f8.jsonl`  
- SHA-256: `86e8241c831035ac6d9dfb0492e5d1c8893245f1e57a2d8432d9926a0cb38de1`  
- Source modified: `2026-07-02T04:49:27+00:00`  
- Imported at: `2026-07-05T16:48:24+00:00`  
- project: `wf_7049e1ed-b93`  
- session_id: `c1fcb4aa-537d-41fd-858e-53f93349bf5a`
```

Transcript

1. user (2026-07-02T04:45:39.413Z)

You are a senior engineer producing ONE isolated PR. Today is 2026-07-02.

REPO: C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector (a shared checkout ? follow the safety rules exactly).

AUDIT FINDINGS YOU ARE FIXING: R1-39 (uncommitted .gitignore hunk) ? Grep

C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-

indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-

guru_2026-07-02.md for these IDs and read the full evidence + verifier notes before coding.

TASK BRIEF:

Resolve the stale uncommitted .gitignore hunk in C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector (an 'artifacts/' addition, ~3 weeks old, self-contradictory vs the ADR it cites ? see the report). Decide with evidence: per ADR-008 / the indexer CODE_MAP, do model-artifact bundles ever live in THIS repo? If yes, commit a CORRECTED pattern (e.g. artifacts/** with ! negations keeping manifests); if no, the right fix is a PR that documents the decision (comment in .gitignore or docs note) WITHOUT the bogus pattern. Either way author the PR from origin/master in your worktree (do NOT reuse the dirty main-checkout hunk ? recreate the correct content). Note in your PR body that the main checkout carries a stale WIP hunk the orchestrator will discard after merge.

BRANCHING SAFETY (mandatory ? shared checkout, sibling sessions exist):

1. In the MAIN repo checkout run: git fetch origin

2. Create an ISOLATED WORKTREE: git worktree add "C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-

indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-<shortname>" -b <branch-name>

origin/master (branch explicitly FROM origin/master ? never from the current local branch; use origin/main for repos whose default is main)

3. Do ALL work inside the worktree. Before committing: git branch --show-current + git log --oneline -1 must show YOUR branch on the origin-master base.

4. Stage EXPLICIT paths only (git add <file> <file>) ? NEVER git add -A / . / -u

5. Before opening the PR: git log --oneline origin/master..HEAD must list ONLY your commits.

COMMIT/PR CONVENTIONS:

- Commit message: conventional style matching the repo's history; end the message with a blank line then: Co-Authoried-By: Claude Fable 5 <noreply@anthropic.com>

- PR body: what/why, a "## Verification" section stating exactly what you ran locally and what CI covers, and end with: ? Generated with [Claude Code](https://claude.com/claude-code)

- Do NOT merge the PR. Do NOT remove the worktree (a fixer may need it). Return the PR number, branch, worktree path.

QUALITY BAR:

- Read the repo's CLAUDE.md first; for the engine also docs/CODE_MAP.md \$0 (file placement) before adding any new file.

- Add/extend tests for every behavior change. Run the targeted tests you touched (python -m pytest tests/<relevant> -q) + ruff check on changed files; if mypy config exists, run mypy on touched modules. Full suite is CI's job ? do not run it.

- Default-safe: new knobs default OFF / behavior-preserving unless the task explicitly says otherwise.

- The full audit evidence is in C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/AUDIT_REPORT_khelsutra-guru_2026-07-02.md ? Grep it for your finding IDs and READ the verifier notes; they contain corrections (exact file:line, scoping traps) you MUST honor.

If the brief's premise turns out false against current code (something already fixed, file moved), adapt honestly and say so in notes ? never force a pointless diff. If you genuinely cannot complete (env broken, auth missing), open NO PR and explain in notes with pr_number=0.

2. user (2026-07-02T04:45:47.287Z)

364-[Omitted long context line]
365-[Omitted long context line]
366-
367-### R1-36 [P2/S] doc-rot ? khelsutra (avidullu/khelsutra) ? SPA web/ + marketing site/
368-**NEXT_STEPS.md was refreshed 2026-07-01 but still advertises the retired droplet host as
'live + browsable'**
369-- anchor: `khelsutra/NEXT\_STEPS.md:12` · tracked in: ? · finder: khelsutra:code
370-[Omitted long context line]
371-- why it matters:
372-- proposed fix: Refresh NEXT_STEPS.md 'Where we are': drop the live-droplet claim, retire
the 'Droplet ? production' queue item (or re-scope it to 'any future VM' ?
site/scripts/provision-vm.sh is already provider-agnostic), and keep the SSH-key revocation as
an explicit still-actionable item. Mark deploy/droplet/README.md as a historical template.
373-[Omitted long context line]
374-[Omitted long context line]
375-
376-### R1-37 [P2/S] doc-rot ? khelsutra (avidullu/khelsutra) ? SPA web/ + marketing site/
377-**All 15 engine line-number contract anchors in the SPA API client are stale, and will drift
further when engine PR #390 (main.py split) merges**
378-- anchor: `khelsutra/web/src/api/client.ts:87` · tracked in: engine PR #390 (the split that
worsens it) · finder: khelsutra:code
379-[Omitted long context line]
380-- why it matters:
381-- proposed fix: Replace hard line numbers with stable route anchors ? e.g. 'backend/main.py
@app.post("/api/process")' / 'backend/api/uploads.py @router.post("/api/upload/init")' ? in one
mechanical pass over client.ts + types.ts. Do it after (or regardless of) #390 so the anchors
survive the god-module split.
382-[Omitted long context line]
383-[Omitted long context line]
384-
385-### R1-38 [P2/M] test-gap ? khelsutra (avidullu/khelsutra) ? SPA web/ + marketing site/
386-**SPA e2e CUJ never runs in CI and the mock engine is hand-synced ? the mock-vs-real drift
bug-class already shipped once**
387-- anchor: `khelsutra/.github/workflows/test.yml:31` · tracked in: e2e/mock-engine.mjs:5
('later drive CI CUJ replays') · finder: khelsutra:code
388-[Omitted long context line]
389-- why it matters:
390-- proposed fix: Add a fourth CI job (mirroring the existing layout job's Playwright setup)
that builds web/ and runs the e2e CUJ specs against mock-engine.mjs; separately add a thin
contract check pinning the mock's response shapes to web/src/api/types.ts (or an engine-side
snapshot of _video_summary/capabilities) so mock-vs-real drift fails a build instead of
shipping.
391-[Omitted long context line]
392-[Omitted long context line]
393-
394-### R1-39 [P2/S] hygiene ? sports-data-collector
395-**Uncommitted .gitignore hunk (artifacts/) stale ~3 weeks and self-contradictory vs the ADR
it cites**
396-- anchor: `.gitignore:59` · tracked in: ? · finder: sdc:audit
397-[Omitted long context line]
398-- why it matters:
399-- proposed fix: Either commit the hunk as a tiny PR with a pattern that actually matches the
stated intent (e.g. 'artifacts/**' + '!artifacts/', '!artifacts/**/',

```
'!artifacts/**/manifest.json'), or discard it if model artifact bundles will only ever live in
the indexer repo (per ADR-008 / indexer CODE_MAP, artifacts/<name>/<ver>/ is an indexer-repo
location). Owner should confirm which repo hosts bundles before committing.
400-[Omitted long context line]
401-[Omitted long context line]
402-
403-### R1-40 [P2/M] ci-infra ? sports-data-collector
404-**CI is pytest-only: no ruff/mypy/coverage gate, and default ruff already fails with 24
errors**
405-- anchor: `.github/workflows/ci.yml:21` · tracked in: ? · finder: sdc:audit
406-[Omitted long context line]
407-- why it matters:
408-- proposed fix: One small PR: (1) fix the 24 ruff findings (20 auto-fixable via --fix), (2)
add [tool.ruff] to pyproject.toml, (3) add 'ruff check .' and 'pytest --cov=src --cov-fail-
under=<current%>' steps to ci.yml. Optionally add a 3.9 job since requires-python >=3.9 is
claimed but only 3.11 is tested.
409-[Omitted long context line]
410-[Omitted long context line]
411-
412-### R1-41 [P2/M] critical-fix ? sports-data-collector
413-**Crawler silently ingests partial matches when a per-set CSV fetch fails (no retry, no
flag)**
414-- anchor: `src/crawlers/badminton_crawler.py:185` · tracked in: docs/NEXT_STEPS.md A1
(retry/backoff only; silent partial-set acceptance not explicitly tracked) · finder: sdc:audit
415-[Omitted long context line]
416-- why it matters:
417-- proposed fix: When any per-set fetch fails, either skip the whole match (mirroring the
rollback-accounting path at :229-237) or record the video with a non-curated status + a
summary_log flag so it is excluded from `curate export`. Retry/backoff convergence onto
download_best_video is the parked A1 item; the silent-partial-accept semantics are NOT
explicitly covered by A1's scope and are the correctness half worth fixing first (pure-local,
testable with a mocked urlretrieve).
418-[Omitted long context line]
419-[Omitted long context line]
420-
421-### R1-42 [P3/S] doc-rot ? badminton-highlight-indexer
422-**deploy/digitalocean/ retired (ADR-013, droplet gone 2026-06-25) but still live-coupled:
carries the trackers pin and an unverified 'rotate the key shared in chat' note**
423-- anchor: `deploy/digitalocean/README.md:3` · tracked in: deploy/digitalocean/README.md:8
(self-noted 'optional later cleanup') · finder: engine:security-config
424-- evidence: deploy/digitalocean/README.md:3-9 marks DO DEPRECATED but says
bootstrap.sh/mount_scratch.sh/gpu_setup.sh are 'still reused by the GCP runbook' and suggests
moving them to deploy/common/; bootstrap.sh:46 is one of the 5 lockstep copies of the trackers
pin; README.md:29 says 'Rotate the old key that was shared in chat' ? rotation is unverifiable
from the repo (hypothesis: may never have happened).
```

3. user (2026-07-02T04:45:49.616Z)

On branch master

Your branch is up to date with 'origin/master'.

Changes not staged for commit:

(use "git add <file>..." to update what will be committed)

(use "git restore <file>..." to discard changes in working directory)

modified: .gitignore

no changes added to commit (use "git add" and/or "git commit -a")

warning: in the working copy of '.gitignore', LF will be replaced by CRLF the next time Git touches it

diff --git a/.gitignore b/.gitignore

index 52dac92..ac1420f 100644

--- a/.gitignore

+++ b/.gitignore

@@ -53,3 +53,7 @@ data/temp/

```
# Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
data/*.backup-*.db
+
+# Model artifact bundles: weights are NEVER committed (private Releases); the
+# reviewable surface is each bundle's manifest.json (ADR-008).
+artifacts/
master
adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)
6f67d65 feat(import): callable import_annotations + a maturity gate (P9) (#52)
c281f82 feat(export): add source file MD5 to the eval/training manifest (P9) (#51)
```

4. user (2026-07-02T04:45:59.779Z)

R1-39 [P2/S] hygiene ? sports-data-collector

```
**Uncommitted .gitignore hunk (artifacts/) stale ~3 weeks and self-contradictory vs the ADR it
cites**
- anchor: `.gitignore:59` · tracked in: ? · finder: sdc:audit
- evidence: git diff shows an uncommitted hunk appending lines 56-59: '# Model artifact bundles:
weights are NEVER committed (private Releases); the reviewable surface is each bundle's
manifest.json (ADR-008).' + 'artifacts/'. File mtime is Jun 10 15:53 ? same day .gitignore was
last committed (e4dd8f1, 2026-06-10) ? so this WIP has sat uncommitted ~3 weeks across sessions.
ADR-008 does not exist in this repo (docs/DECISIONS.md has only ADR-001/ADR-002); it lives in
the indexer at badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md:388 ('Owned-
model topology'). Internal contradiction: a blanket 'artifacts/' ignore also hides the
manifest.json the comment declares reviewable (same contradiction exists in the indexer: its
.gitignore:53 ignores artifacts/ wholesale while its CODE_MAP.md:35 claims 'manifest
committed'). No artifacts/ dir exists in this repo yet.
- why it matters:
- proposed fix: Either commit the hunk as a tiny PR with a pattern that actually matches the
stated intent (e.g. 'artifacts/**' + '!artifacts/', '!artifacts/**/',
'!artifacts/**/manifest.json'), or discard it if model artifact bundles will only ever live in
the indexer repo (per ADR-008 / indexer CODE_MAP, artifacts/<name>/<ver>/ is an indexer-repo
location). Owner should confirm which repo hosts bundles before committing.
- verifier notes: Sharper failure scenario: a future session in this shared checkout commits
.gitignore (or an agent stages broadly) and the hunk lands unreviewed, encoding a self-
contradictory policy citing an ADR unresolvable in-repo; later, if artifact bundles are ever
produced here, manifest.json files are silently invisible to `git status`, so the declared
provenance/review surface never gets committed ? the same silent gap already real in the indexer
(CODE_MAP.md:35 says "manifest committed" but `git ls-files artifacts/` is empty). Minor
correction: the indexer .gitignore's own comment at line 52 says artifacts are "regenerable,
never committed", so the CODE_MAP-vs-gitignore contradiction there is doc-vs-doc, not just
pattern-vs-comment. Severity P2 stands as a hygiene finding given the documented prior squash-
merge absorption incident in this multi-session checkout; effort S stands (one tiny PR or a one-
line `git checkout -- .gitignore` after owner direction).
- verifier reasoning: Reproduced every element of the claim in the current tree. (1) `git diff
-- .gitignore` in sports-data-collector shows exactly the claimed uncommitted hunk appended
after line 55 (`data/*.backup-*.db`): a two-line comment "Model artifact bundles: weights are
NEVER committed (private Releases); the reviewable surface is each bundle's manifest.json
(ADR-008)." plus `artifacts/` at working-tree line 59; `git status --porcelain` shows ` M
.gitignore` on master at adcd206. (2) Staleness confirmed: last commit touching .gitignore is
e4dd8f1 dated 2026-06-10 14:42:02 +0530, file mtime 2026-06-10 15:53, today 2026-07-02 ? 22 days
uncommitted while master advanced through PRs #51?#53. Not fixed by any later commit (git log --
.gitignore shows nothing after e4dd8f1). (3) ADR-008 absent from this repo: sports-data-
collector/docs/DECISIONS.md contains only ADR-001 (line 8) and ADR-002 (line 76); ADR-008 lives
at badminton-highlight-indexer/docs/archives/decisions/DECISIONS.md:388 ("ADR-008: Owned-model
topology"). (4) Internal contradiction confirmed: a blanket `artifacts/` ignore hides
manifest.json too, so the comment's "reviewable surface is manifest.json" is unenforceable as
written. Same contradiction verified in the indexer: badminton-highlight-indexer/.gitignore:53
ignores `artifacts/` wholesale (comment at :52 even says "never committed"), while
docs/CODE_MAP.md:35 claims "manifest committed (provenance), weights gitignored" ? and `git ls-
files artifacts/` in the indexer returns zero tracked files, i.e. no manifest is actually
committed anywhere. (5) No artifacts/ dir exists in sports-data-collector ? impact is latent,
but the hygiene hazard is live: a stale WIP hunk in a checkout shared by concurrent sessions can
ride into an unrelated commit (the owner's own global conventions document a prior incident of
exactly this). Proposed fix is sane and correctly scoped: the un-ignore recipe (`artifacts/**` +
```

`!artifacts/\*\*/` + `!artifacts/\*\*/manifest.json`) is valid gitignore semantics (`!artifacts/` is redundant but harmless), and the discard alternative matches CODE_MAP.md:35 placing artifacts/<name>/<ver>/ in the indexer repo. Not owner-gated per the listed gates (no serving go-live, corpus promotion, licensing, spend, or counsel involvement) ? it only needs the owner's one-line direction call on which repo hosts bundles, which the proposed fix already states.

5. user (2026-07-02T04:46:00.150Z)

File does not exist. Note: your current working directory is C:\Users\avidu\Projects\kheIsutra-guru\badminton-highlight-indexer.

6. user (2026-07-02T04:46:06.630Z)

No files found

7. user (2026-07-02T04:46:07.834Z)

README.md
__pycache__
cleanup.py
config
data
docs
pyproject.toml
src
tests
DECISIONS.md
INGESTION_PIPELINE.md
MANUAL_INGESTION_POSITIONING.md
NEXT_STEPS.md
PIPELINE_ROLE.md
annotation
archives

8. user (2026-07-02T04:46:22.400Z)

Already up to date.

master...origin/master

M .gitignore

9. user (2026-07-02T04:46:22.569Z)

WASB's candidate windows and the *licensing question* on its academic-trained weights.

****Why we don't jump to Option B now ? and why Option A is the path to it.**** Option B needs data we don't yet have, and it's far harder to debug as one black box. So we deliberately build ****modular cues first**** ? audio (this ADR), then a stance heuristic (the parked #2), then other signals ? each cheap, debuggable, label-efficient, and *each one generating the very golden labels an end-to-end model will need*. When the corpus is healthy, we train the unified model. ****Option A is not a detour from Option B; it is how we earn it.****

ADR-008: Owned-model topology ? artifact boundary now, repo split at Gen-2

****Date:**** 2026-06-10

****Status:**** Accepted (owner-ratified after an adversarial design panel)

Context

M0.4 (Gen-0 TAL train?eval?ship-gate harness, productionizing `backend/eval/rally\_seq\_proto.py`

over ActionFormer/ASFormer) and M0.1 (corpus spine + event-index contract) are greenlit. The question: should the model/training work live **in this repo**, or in a **separate repo** whose output the indexer consumes as another segmenter (like WASB and Gemini)? The owner's instinct was isolation ("iterate on model and processing independently"), citing the collector split as precedent.

Decision

Draw the boundary at the ARTIFACT CONTRACT now; split the repo at Gen-2, as a pre-committed, trigger-based event ? not today.

- Training lives in this repo as a top-level `training/` tree** (sibling of `backend/`), isolated **executably**, not socially: its own `training/requirements-train.txt` (heavy/training-only deps never enter `requirements.txt`), and the import-smoke suite asserts ``import backend.main`` loads **no `training.\*` module**.
- The featurizer is single-source-of-truth**: promoted out of `rally\_seq\_proto` into `backend/features/rally\_seq.py` (pure numpy/scipy, versioned schema constant), imported by BOTH the training harness and the future TAL serving segmenter. Train == serve by construction ? the fps-normalization bug (fixed 2026-06-10) is the skew class this kills.
- Serving consumes only a versioned ARTIFACT BUNDLE** (weights gitignored / private GitHub Releases + a manifest carrying: feature-schema version, normalization, calibration/post-processing, training-corpus ref + GT hash, consent + license attestation incl. the WASB C3 checkpoint status). The loader/gate hard-fails on schema mismatch ? the `regression.py` incommensurability-guard pattern.
- The ship-gate is owned by the product side** (extends `backend/eval/regression.py` + `metrics.match\_intervals`): the trainer runs it as pre-flight but never grades itself. The gate refuses commercial ship while C3 is unresolved, regardless of F1.
- Pre-committed split trigger** ? extract `training/` into a separate repo (working name `sports-temporal-models`, to be finalized with the pending repo rename) **when**:
(a) Gen-2 arrives (ms-swift / Qwen3-VL / cloud-GPU training ? an alien dependency tree), OR
(b) a second consumer of the trained artifacts appears (e.g. the khelacharya layer), AND (c) the harness's only data input is the collector's M0.1 export. By then the artifact contract is proven and extraction is a file move across a known seam.

Prerequisite frozen first: the golden-label/corpus contract. One canonical GT loader (the repo carried two incompatible ones: ``annotations.load_annotations`` `start,end` vs ``calibrate_wasb.load_golden_gts`` `start\_time,end\_time`, the latter feeding 7 modules) + a round-trip test against the collector's ``curate export`` format. Every ship-gate quantity (gt_hash, LOVO numbers, the future consent fingerprint) flows through this boundary.

Rationale (the decisive, code-verified facts)

- The "just another runner like WASB/Gemini" analogy fails at the seam location.** WASB's contract sits at the video boundary (video ? trajectory CSV); Gemini's at video ? segments. The Gen-0 head consumes **features this repo computes** from the trajectory ? the seam is

10. user (2026-07-02T04:46:34.747Z)

```
| You're adding? | Put it in? | Notes |
|---|---|---|
| Backend / runtime code | `backend/<subsystem>/` |
`pipeline/{segmenters,detectors,validators}`, `providers/`, `sports/`, `annotations/`,
`infrastructure/`, `reporting/`, `stitcher/`, `utils/`, `config/`, `features/`, `api/`. Most
additions are register a class in the relevant `?` registry, not a new module tree. |
| Eval LIBRARY (importable, pure ? metrics, calibration, a scorer, a feature) |
`backend/eval/` | Pure core, no side-effects at import; CLI/I-O lives in a separate driver.
These are imported widely (e.g. `calibrate_wasb` by 8 sites) ? they are library, not
scripts; do NOT move them out. |
| Standalone run-driver / one-off entry CLI (has `__main__`, NOT imported by app code) |
scripts/` | e.g. `scripts/run_phase3_eval.py`, `scripts/run_process_and_stitch.py`. Run as
`python scripts/<name>.py`. If it needs `load_dotenv()` / `sys.path` before importing backend, add
`"scripts/<name>.py" = ["E402"]` to `ruff.toml`. |
| Ops / maintenance tool (not part of the served app ? cache GC, batch admin) | `tools/` |
Run as `python -m tools.<name>` (it's an importable package). e.g. `tools/cleanup_caches.py`.
```

Keep the destructive decision path pure + unit-tested. |
 | Training code (model fitting) | `training/` | Behind the CI-enforced import wall (ADR-008) ?
 `backend.main` must not import it. Own `requirements-train.txt`. |
A test	`tests/test\_\*.py`	pytest auto-discovers; the full-suite lane has a **70% coverage floor**. Pure/offline + mocked (no GPU/WSL/network).
A doc (design, plan, living log)	`docs/`	Index it in `docs/README.md`. Only **terminal-state** (DONE/REJECTED) work moves to `docs/archives/`.
A throwaway repro / debug script	`scratch/`	**Gitignored** + excluded from ruff/pytest. Never imported by app code.
A model artifact (weights + manifest)	`artifacts/<name>/<ver>/`	ADR-008: **manifest committed** (provenance), **weights gitignored**.
Eval baseline / leaderboard JSON	`eval\_baselines/`	Tracked (survives a clean checkout); the no-regression gate reads it.
Committed regression fixture (video-free golden)	`eval\_baselines/fixtures/<slug>/`	Tracked, **LF-pinned** (`.gitattributes`). Synthetic Tier-0 (or owner-cleared, de-attributed real). Drives the `golden\_fixtures` characterization gate with **no video/GPU/DB**. See `docs/GOLDEN\_REGRESSION\_FIXTURES.md`; mint via `python -m backend.eval.golden\_fixtures --freeze-synthetic` / `--freeze-real` (refusal-gated).
Pipeline run outputs (reels, trajectories, DBs, frames)	`output/`	**Gitignored**. Per-video isolation is the tracked **PER_VIDEO_WORKSPACE.md** project (P0 helper shipped #264; supersedes the archived `PER\_VIDEO\_OUTPUT\_LAYOUT`).

Canonical files that stay at the repo ROOT (referenced by CI / packaging / docs ? do not move):
 `pyproject.toml` (PEP 621 packaging, hatchling; `indexer` entry point) · `run\_all\_tests.py` (CI: `ci.yml`) ·
 `run\_regression.py` (the no-regression gate CLI, cited in `EXPERIMENT\_HARNESS.md`) ·
 `run\_eval.py`
 (grandfathered ? imported by `tests/test\_eval\_harness.py`). The old diagnostic `setup.py` moved to
 `scripts/doctor.py` (it shadowed the build system); invoke via `python scripts/doctor.py` or
 `indexer --setup`.
 New standalone run-drivers go in **scripts/`**, not the root.

11. user (2026-07-02T04:46:35.671Z)

```
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\.gitignore:46:# ---- Generated temp
artifacts (crawler scratch space; regenerated each run) ----
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\.gitignore:57:# Model artifact
bundles: weights are NEVER committed (private Releases); the
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\.gitignore:59:artifacts/
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\docs\DECISIONS.md:130: precedent
for the failure mode: unfilled intake-manifest TODOs, eval artifacts lost in
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\tests\test_cvat_robustness.py:5:uncaught ParseError; the standard .zip delivery
artifact was unsupported.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\tests\test_cvat_parser.py:158:    #
and they survive a reload (the CSV is the durable artifact)
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\tests\test_cvat_parser.py:349:# ---
validate-delivery: gap heuristic, CSV loading, report artifacts -----
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\tests\test_cvat_parser.py:391:def
test_validation_artifacts(tmp_path):
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\src\reporting\__init__.py:38:
<id>_report.* artifacts). Used by validate-delivery / convert-cvat. Never raises.""
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\src\cli\curate.py:826:            # or
unverified file under the to-be-shared name is exactly the artifact
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\src\cli\curate.py:1111:        #
colocated with the existing artifacts. Never raises.
C:\Users\avidu\Projects\khelsutra-guru\sports-data-collector\src\parsers\cvat\rally_cvat.py:231:
# Accept the standard CVAT delivery artifact (a .zip containing annotations.xml) as
C:\Users\avidu\Projects\khelsutra-guru\sports-data-
collector\src\parsers\cvat\rally_cvat.py:590:# These turn a (records, issues) pair into the
artifacts a human reviews:    #
```

12. assistant (2026-07-02T04:46:56.536Z)

The evidence is decisive. Now let me look at the committed .gitignore content (at origin/master) and the repo's commit history style.

13. user (2026-07-02T04:46:56.886Z)

Exit code 128

fatal: ambiguous argument 'origin\master;.gitignore': unknown revision or path not in the working tree.

Use '--' to separate paths from revisions, like this:

'git <command> [<revision>...] -- [<file>...]'

14. user (2026-07-02T04:46:57.779Z)

adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)
6f67d65 feat(import): callable import_annotations + a maturity gate (P9) (#52)
c281f82 feat(export): add source file MD5 to the eval/training manifest (P9) (#51)
a152cf5 fix(cvat): warn on rows silently skipped by load_canonical_csv (#50)
75bb0fa docs: archive CODE_HEALTH_PLAN ? project DONE (#49)
a09a000 refactor(manifests): typed pydantic models for the export manifest (A5) (#48)
6678fb0 refactor(curate): lift sys.exit out of _guard_import into CurationError (A2) (#47)
9e56056 fix(reporting): make FAQ path configurable via env var (A6) (#46)
0ddb022 docs: Tier D triage (D8) + NEXT_STEPS for parked A1/A3 (#45)
c7c4f68 refactor(config): promote validation thresholds to config (Tier C) (#43)
288b8d7 refactor(curate): concrete type hints for annotation handling (Tier C) (#44)
049cf69 refactor(db): versioned schema migrations via PRAGMA user_version (T3) (#41)
9f1e792 fix(rally_cvat): harden XML parsing with defusedxml (T7/QQ5) (#42)
1706180 refactor(schemas): extract shared timestamp validator to mixin (T2) (#40)
768968e refactor(validator,parsing): dedupe license whitelist + strict safe_int (T9/T6) (#39)
origin https://github.com/Khelsutra/sports-data-collector.git (fetch)
origin https://github.com/Khelsutra/sports-data-collector.git (push)
github.com
? Logged in to github.com account avidullu (keyring)
- Active account: true
- Git operations protocol: https
- Token: gho_*****

15. user (2026-07-02T04:47:06.539Z)

Exit code 128

fatal: ambiguous argument 'origin\master;.gitignore': unknown revision or path not in the working tree.

Use '--' to separate paths from revisions, like this:

'git <command> [<revision>...] -- [<file>...]'

16. user (2026-07-02T04:47:14.417Z)

---- Python ----

__pycache__/
*.py[cod]
*\$py.class
*.egg-info/
.eggs/
build/
dist/
.installed.cfg
*.egg

---- Test / coverage ----

.pytest_cache/
.coverage
.coverage.*
htmlcov/
.tox/

.cache/

---- Virtual environments ----

.venv/
venv/
env/
ENV/

---- Secrets / local config ----

.env
.env.*

---- OS / editor cruft ----

.DS_Store
Thumbs.db
.idea/
.vscode/
*.swp

---- Project: large media kept outside git (see README) ----

The video corpus is hundreds of MB+; only metadata (the SQLite DB) and the small source annotation CSVs are version-controlled. Re-fetch videos from the stored video_url / local_path per the manifest.

data/videos/

---- Generated eval/training exports (regenerated via `curate export`) ----

data/eval_export*/

---- Generated temp artifacts (crawler scratch space; regenerated each run) ----

data/temp/

---- Per-video observability reports (written next to videos / inputs at runtime) ----

*.report.json
*.report.md
*.summary.md

Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit

data/*.backup-*.db

17. user (2026-07-02T04:47:44.822Z)

Preparing worktree (new branch 'docs/gitignore-artifact-bundles-decision')
branch 'docs/gitignore-artifact-bundles-decision' set up to track 'origin/master'.
HEAD is now at adcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)

18. user (2026-07-02T04:47:51.047Z)

```
50      *.report.json
51      *.report.md
52      *.summary.md
53
54      # Local DB backups (e.g. data/sports_metadata.backup-YYYY-MM-DD.db) ? never commit
55      data/*.backup-*.db
56
```

19. user (2026-07-02T04:48:00.715Z)

The file C:/Users/avidu/AppData/Local/Temp/claude/C--Users-avidu-Projects-badminton-highlight-indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-gitignore-artifacts/.gitignore has been updated successfully. (file state is current in your context ? no need to Read it back)

20. assistant (2026-07-02T04:48:05.875Z)

Now verify the .gitignore is functionally unchanged (comment-only) and existing patterns still work.

21. user (2026-07-02T04:48:09.416Z)

```
--- should be ignored (existing patterns):
.gitignore:41:data/videos/      data/videos/x.mp4
.gitignore:47:data/temp/       data/temp/t.csv
.gitignore:55:data/*.backup-*.db data/sports_metadata.backup-2026-01-01.db
.gitignore:50:*.report.json    foo.report.json
--- should NOT be ignored (exit 1 expected):
exit=1
```

22. user (2026-07-02T04:48:18.509Z)

```
docs/gitignore-artifact-bundles-decision
addcd206 fix(db): set SQLite busy_timeout so concurrent writers wait, not error (P9) (#53)
M .gitignore
```

23. user (2026-07-02T04:48:28.341Z)

```
[docs/gitignore-artifact-bundles-decision ab6f68b] docs gitignore: record that model-artifact
bundles never live in this repo (ADR-008)
 1 file changed, 14 insertions(+)
ab6f68b docs gitignore: record that model-artifact bundles never live in this repo (ADR-008)
```

24. assistant (2026-07-02T04:48:33.877Z)

Branch contains only my commit. Push and open the PR.

25. user (2026-07-02T04:48:38.515Z)

```
remote:
remote: Create a pull request for 'docs/gitignore-artifact-bundles-decision' on GitHub by
visiting:
remote:   https://github.com/Khelsutra/sports-data-collector/pull/new/docs/gitignore-
artifact-bundles-decision
remote:
branch 'docs/gitignore-artifact-bundles-decision' set up to track 'origin/docs/gitignore-
artifact-bundles-decision'.
To https://github.com/Khelsutra/sports-data-collector.git
 * [new branch]      docs/gitignore-artifact-bundles-decision -> docs/gitignore-artifact-
bundles-decision
```

26. user (2026-07-02T04:49:00.609Z)

<https://github.com/Khelsutra/sports-data-collector/pull/54>

27. assistant (2026-07-02T04:49:08.334Z)

PR opened. Final state check on the main checkout to confirm it was left untouched.

28. user (2026-07-02T04:49:09.504Z)

```
M .gitignore
```

```
C:/Users/avidu/Projects/khelsutra-guru/sports-data-collector
adcd206 [master]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-curate-verify      adcd206
[fix/curate-verify-registry-rot]
C:/Users/avidu/AppData/Local/Temp/claude/c--Users-avidu-Projects-badminton-highlight-
indexer/c1fcb4aa-537d-41fd-858e-53f93349bf5a/scratchpad/wt-gitignore-artifacts  ab6f68b
[docs/gitignore-artifact-bundles-decision]
```

29. user (2026-07-02T04:49:27.590Z)

Structured output provided successfully